

(12) **United States Patent**
Gee et al.

(10) **Patent No.:** **US 12,393,690 B2**
(45) **Date of Patent:** **Aug. 19, 2025**

(54) **RECOVERING QUARANTINED
INFORMATION FROM BACKUP
LOCATIONS**

(71) Applicant: **Rubrik, Inc.**, Palo Alto, CA (US)

(72) Inventors: **Adam Gee**, San Francisco, CA (US);
Kunal Sean Munshani, Fremont, CA
(US); **Benjamin Travis Meadowcroft**,
Lincoln, CA (US); **Guilherme Vale**
Ferreira Menezes, Los Gatos, CA
(US); **Karthick Raja Ravichandran**,
Sunnyvale, CA (US); **William Michael**
Davis, Brisbane (AU); **Andrew**
William Draper, Newcastle-upon-Tyne
(GB); **Haijin He**, South Setauket, NY
(US); **Xiaoyang Zhao**, Fremont, CA
(US)

(73) Assignee: **Rubrik, Inc.**, Palo Alto, CA (US)

(*) Notice: Subject to any disclaimer, the term of this
patent is extended or adjusted under 35
U.S.C. 154(b) by 330 days.

(21) Appl. No.: **17/980,752**

(22) Filed: **Nov. 4, 2022**

(65) **Prior Publication Data**

US 2023/0177158 A1 Jun. 8, 2023

Related U.S. Application Data

(60) Provisional application No. 63/421,536, filed on Nov.
1, 2022, provisional application No. 63/319,953, filed
(Continued)

(51) **Int. Cl.**
G06F 21/56 (2013.01)

(52) **U.S. Cl.**
CPC **G06F 21/568** (2013.01)

(58) **Field of Classification Search**
CPC G06F 21/568
See application file for complete search history.

(56) **References Cited**

U.S. PATENT DOCUMENTS

8,495,037 B1 7/2013 Westenberg
9,218,252 B1 12/2015 Revur et al.

(Continued)

OTHER PUBLICATIONS

ISA/EP, Int'l App. No. PCT/US2022/079400, Partial International
Search Report and Provisional Opinion dated Mar. 7, 2023, 6 pages.

(Continued)


Primary Examiner — Clayton R Williams

(74) *Attorney, Agent, or Firm* — Holland & Hart LLP

(57) **ABSTRACT**

Subject matter related to data management is disclosed. Snapshots of a target object in a computing system may be stored, where the snapshot may correspond to multiple versions of the target object. An indication that information in one or more of the snapshots may be stored. A request to restore a snapshot of the target object may be received. In response to the request, a determination of whether the snapshot of the target object is prohibited from being restored may be made based on the indication that information in one or more of the snapshots is flagged for quarantine. A determination of whether to initiate a restoration operation in response to the request to restore the snapshot may also be made based on whether the snapshot is prohibited from being restored.

20 Claims, 16 Drawing Sheets



Related U.S. Application Data

on Mar. 15, 2022, provisional application No. 63/276,822, filed on Nov. 8, 2021.

2021/0240828	A1	8/2021	Gaurav et al.
2022/0100378	A1	3/2022	Borate et al.
2022/0245250	A1	8/2022	Warwick et al.
2022/0345473	A1	10/2022	Kare et al.

(56)

References Cited

U.S. PATENT DOCUMENTS

10,650,146	B1 *	5/2020	Gaurav	G06F 21/562
10,887,339	B1 *	1/2021	Sokolov	H04L 63/145
11,681,591	B2 *	6/2023	Kulaga	G06F 11/3006
				726/23
2007/0245105	A1	10/2007	Suzuki et al.	
2011/0197279	A1 *	8/2011	Ueoka	G06F 21/568
				726/24
2019/0235973	A1	8/2019	Brewer et al.	
2019/0354443	A1 *	11/2019	Haustein	G06F 11/1451
2020/0159624	A1	5/2020	Malkov et al.	
2020/0201998	A1	6/2020	Jung	
2020/0226256	A1	7/2020	Gaurav et al.	
2020/0319979	A1 *	10/2020	Kulaga	G06F 11/3006
2021/0044604	A1	2/2021	Annen et al.	

OTHER PUBLICATIONS

ISA/EP, Int'l App. No. PCT/US2022/079400, Search Report and Written Opinion dated Mar. 7, 2023, 9 pages.
 U.S. Appl. No. 17/980,645, filed Nov. 4, 2022, Pending, Indicating Infected Snapshots in a Snapshot Chain.
 U.S. Appl. No. 17/980,652, filed Nov. 4, 2022, Pending, Recovering Infected Snapshots in a Snapshot Chain.
 U.S. Appl. No. 17/980,676, filed Nov. 4, 2022, Pending, Quarantining Information in Backup Locations.
 U.S. Appl. No. 17/980,930, filed Nov. 4, 2022, Pending, Bulk Snapshot Recovery.
 PCT/US22/79400, Nov. 7, 2022, Pending, Snapshot-Based Malware Management.

* cited by examiner




FIG. 1




FIG. 2

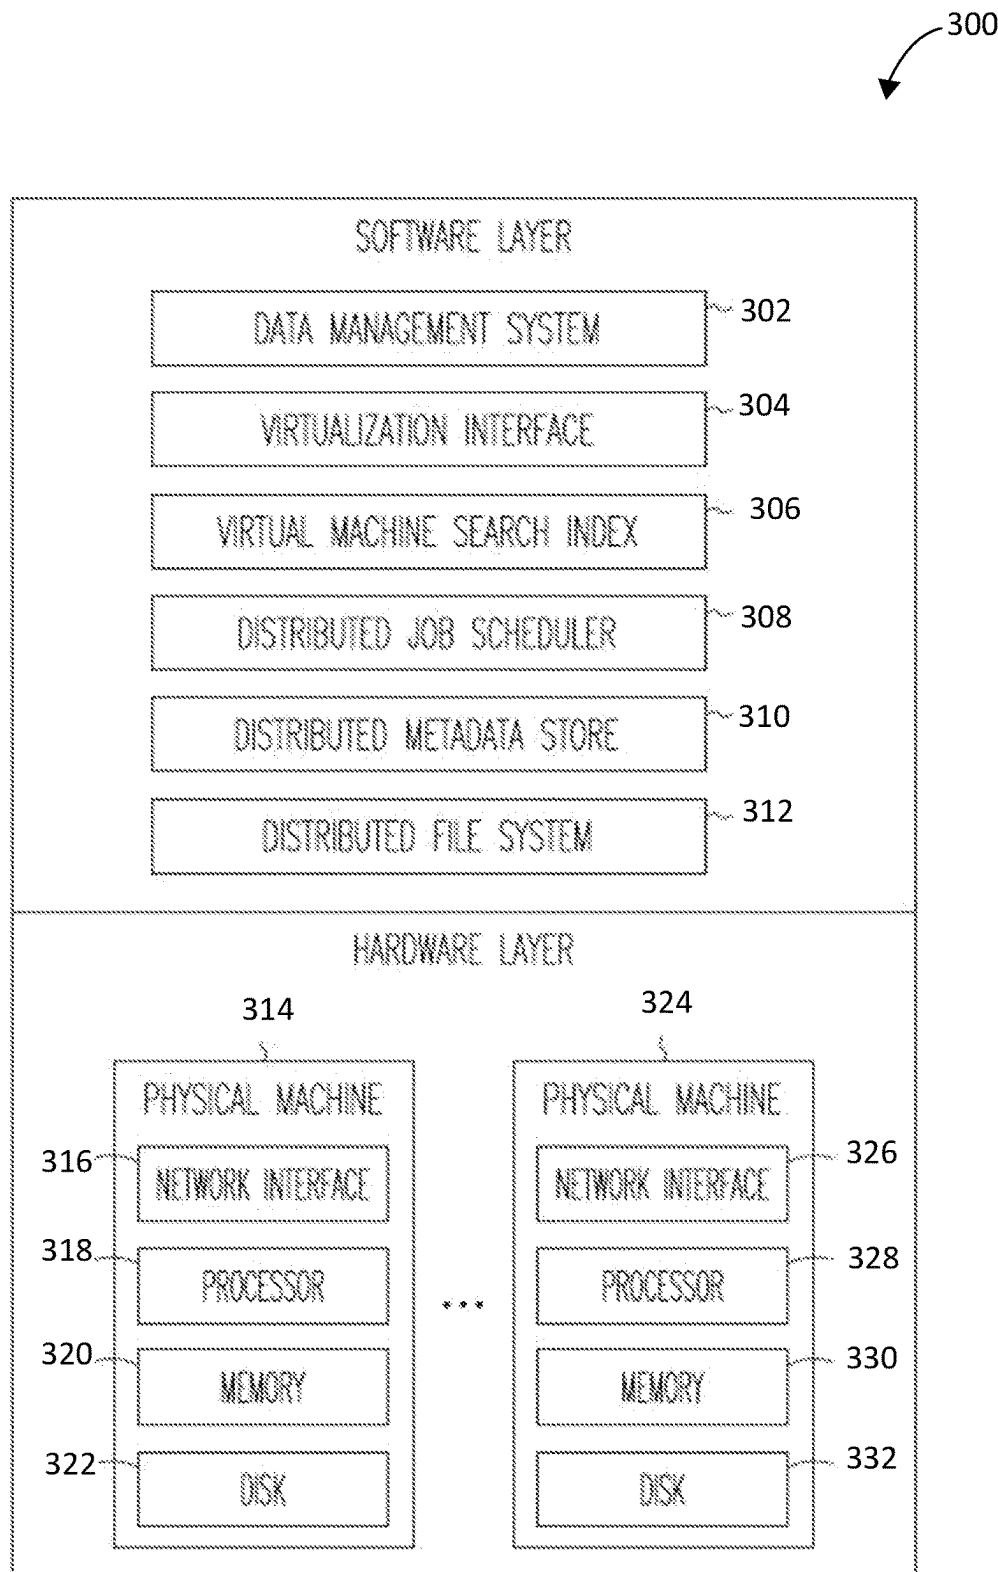


FIG. 3




FIG. 4

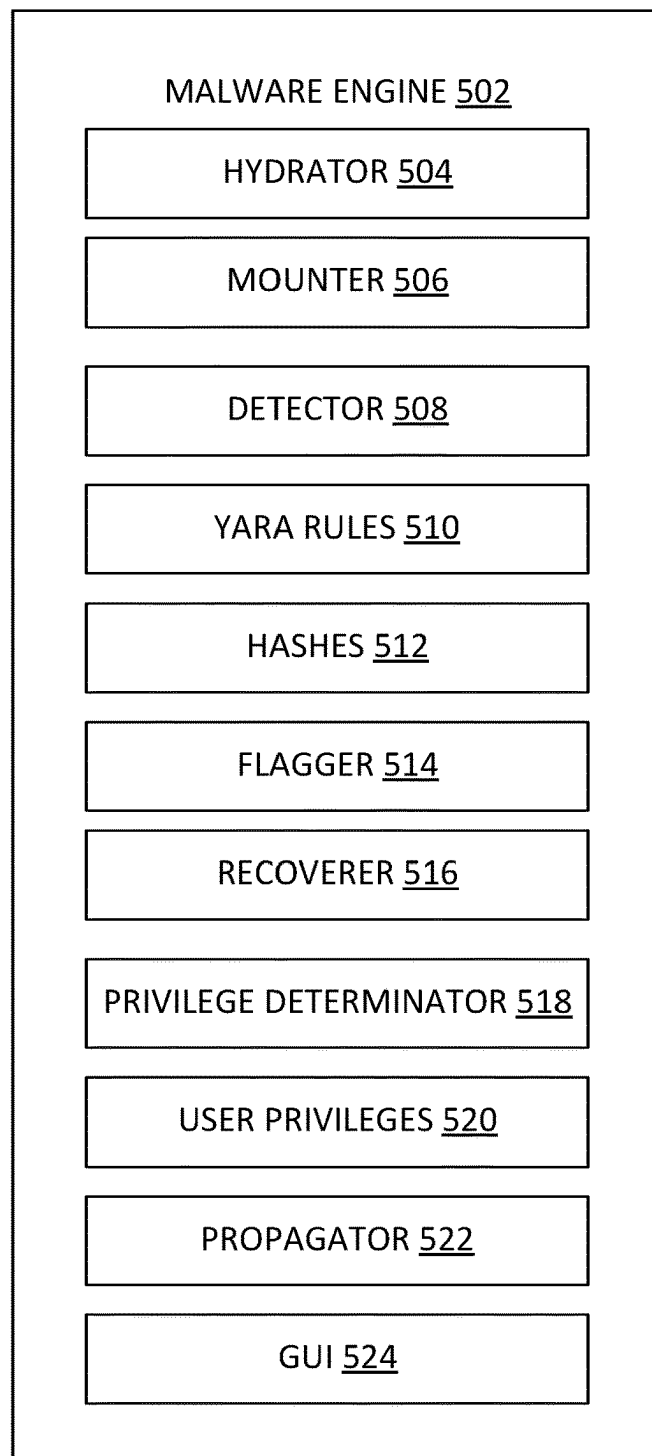


FIG. 5

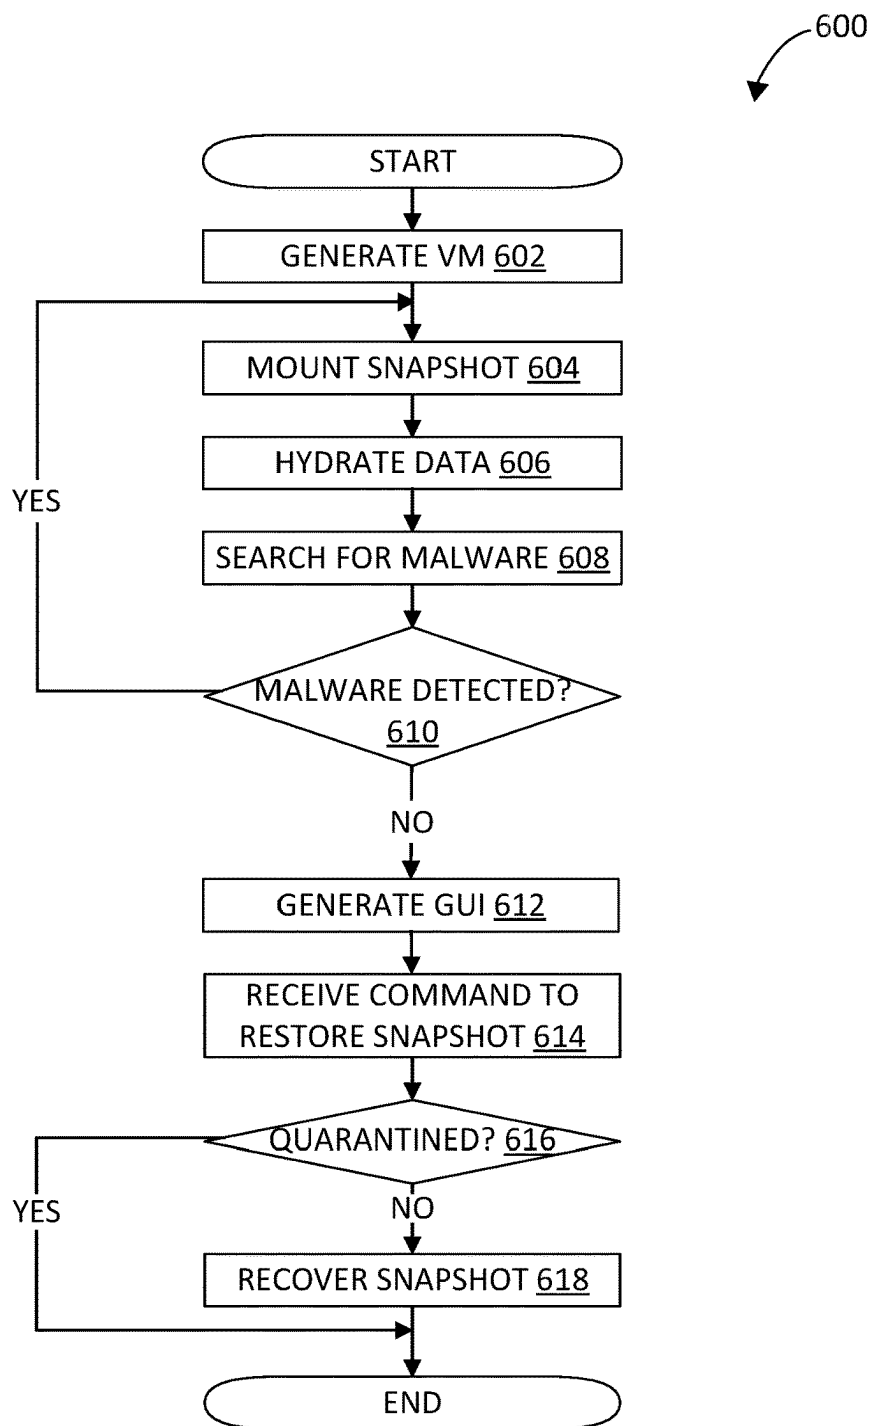


FIG. 6

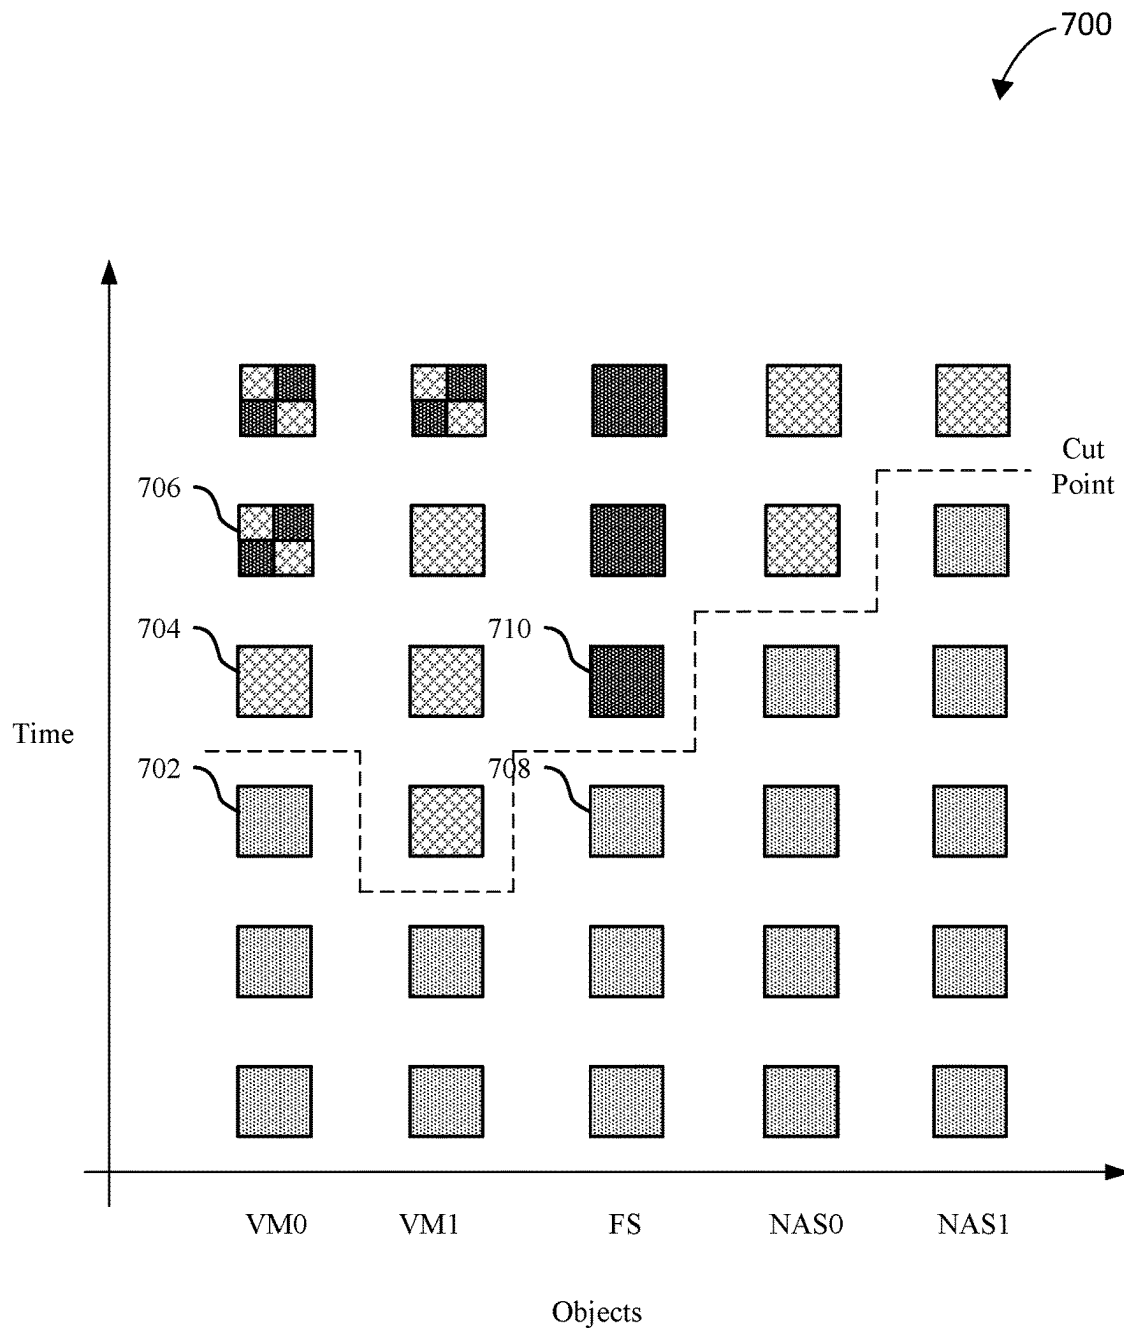


FIG. 7

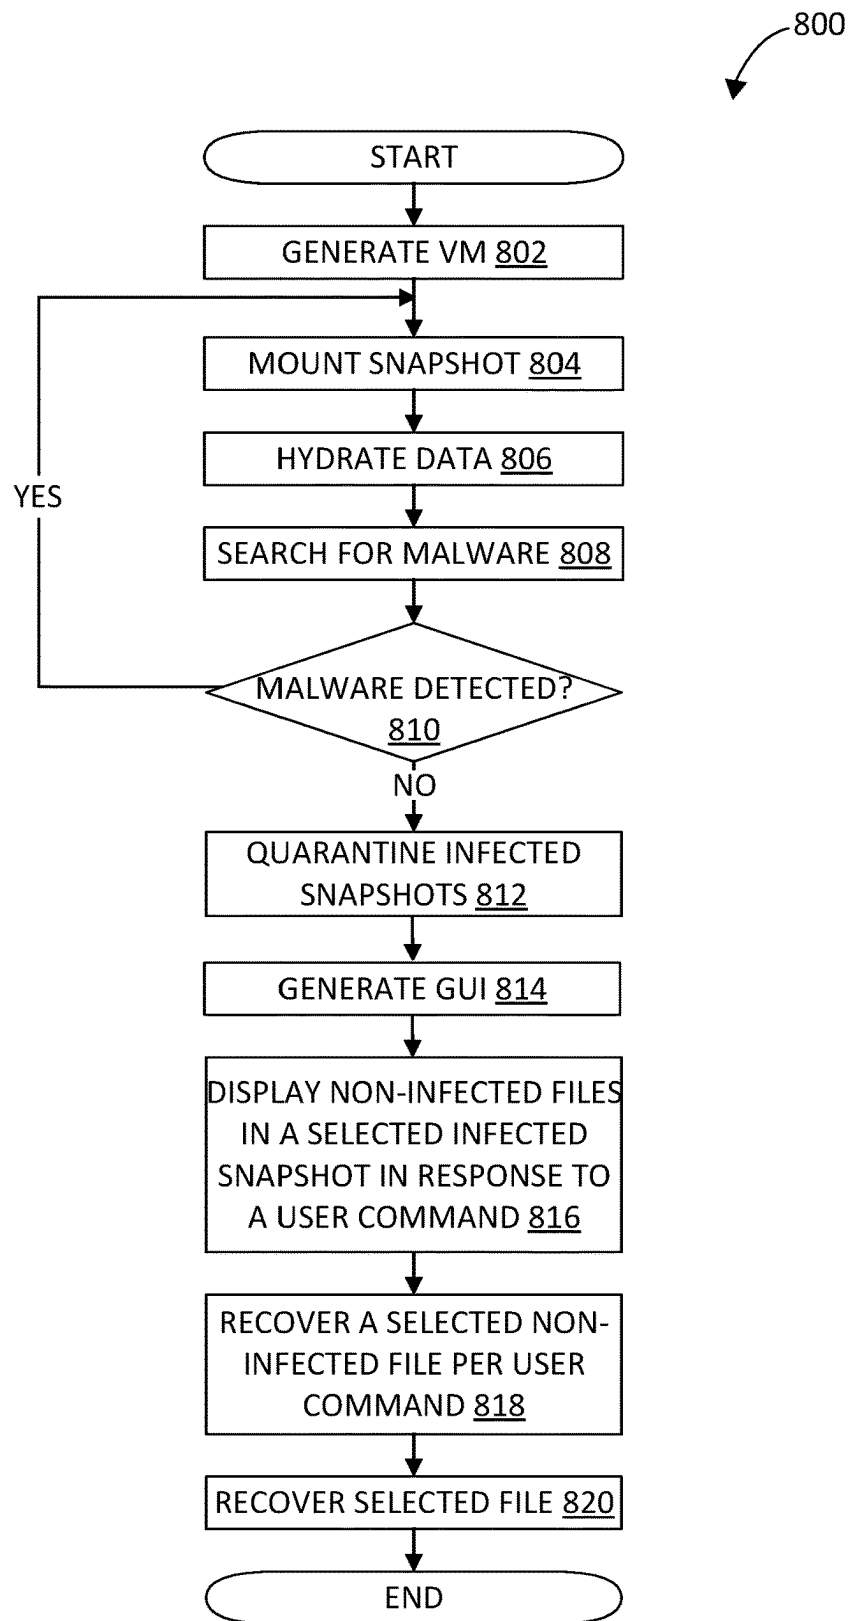


FIG. 8

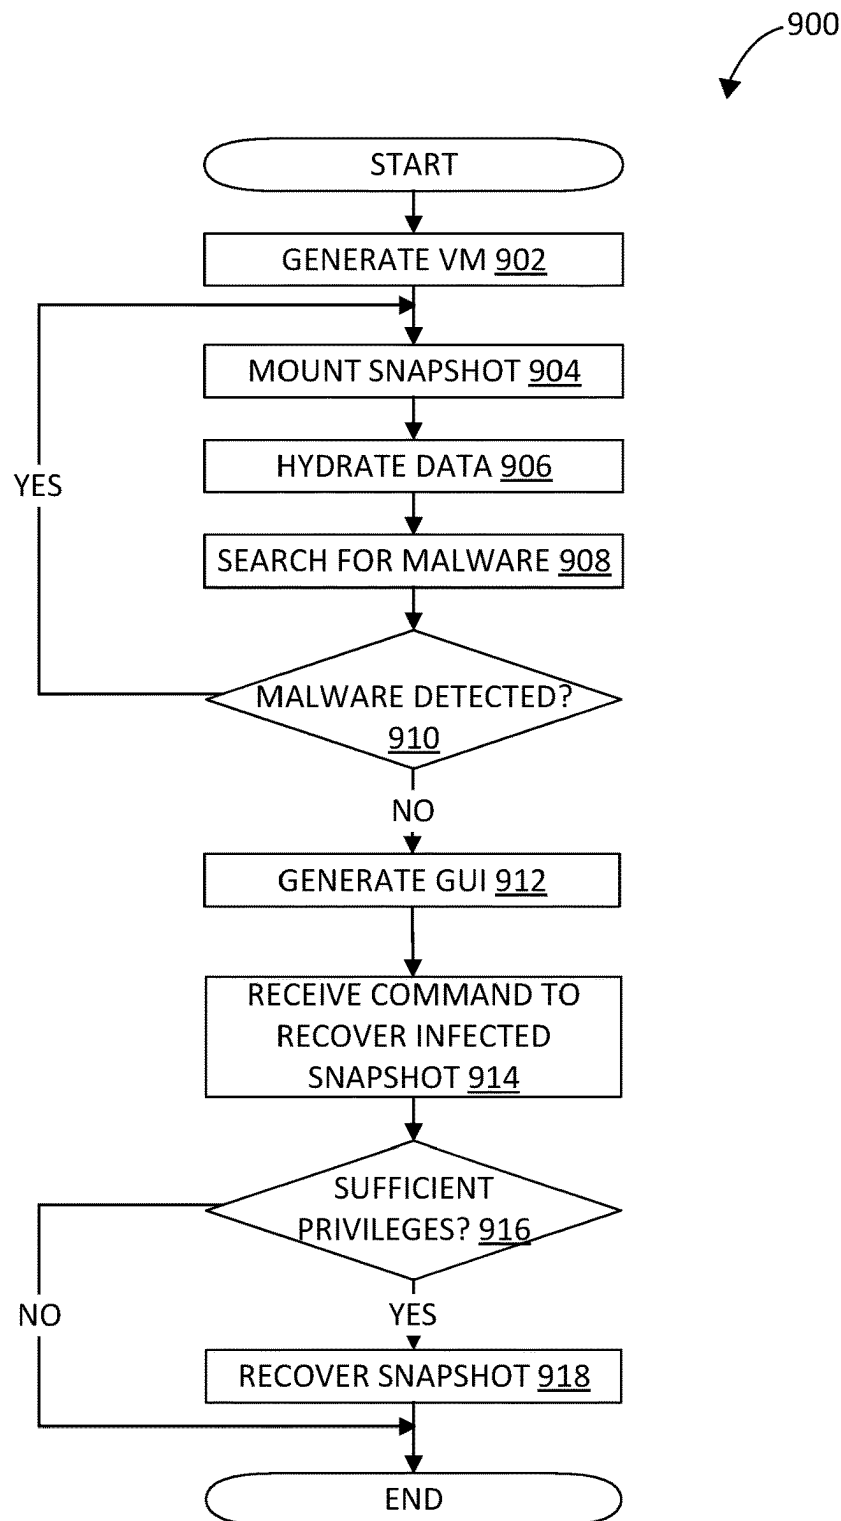


FIG. 9

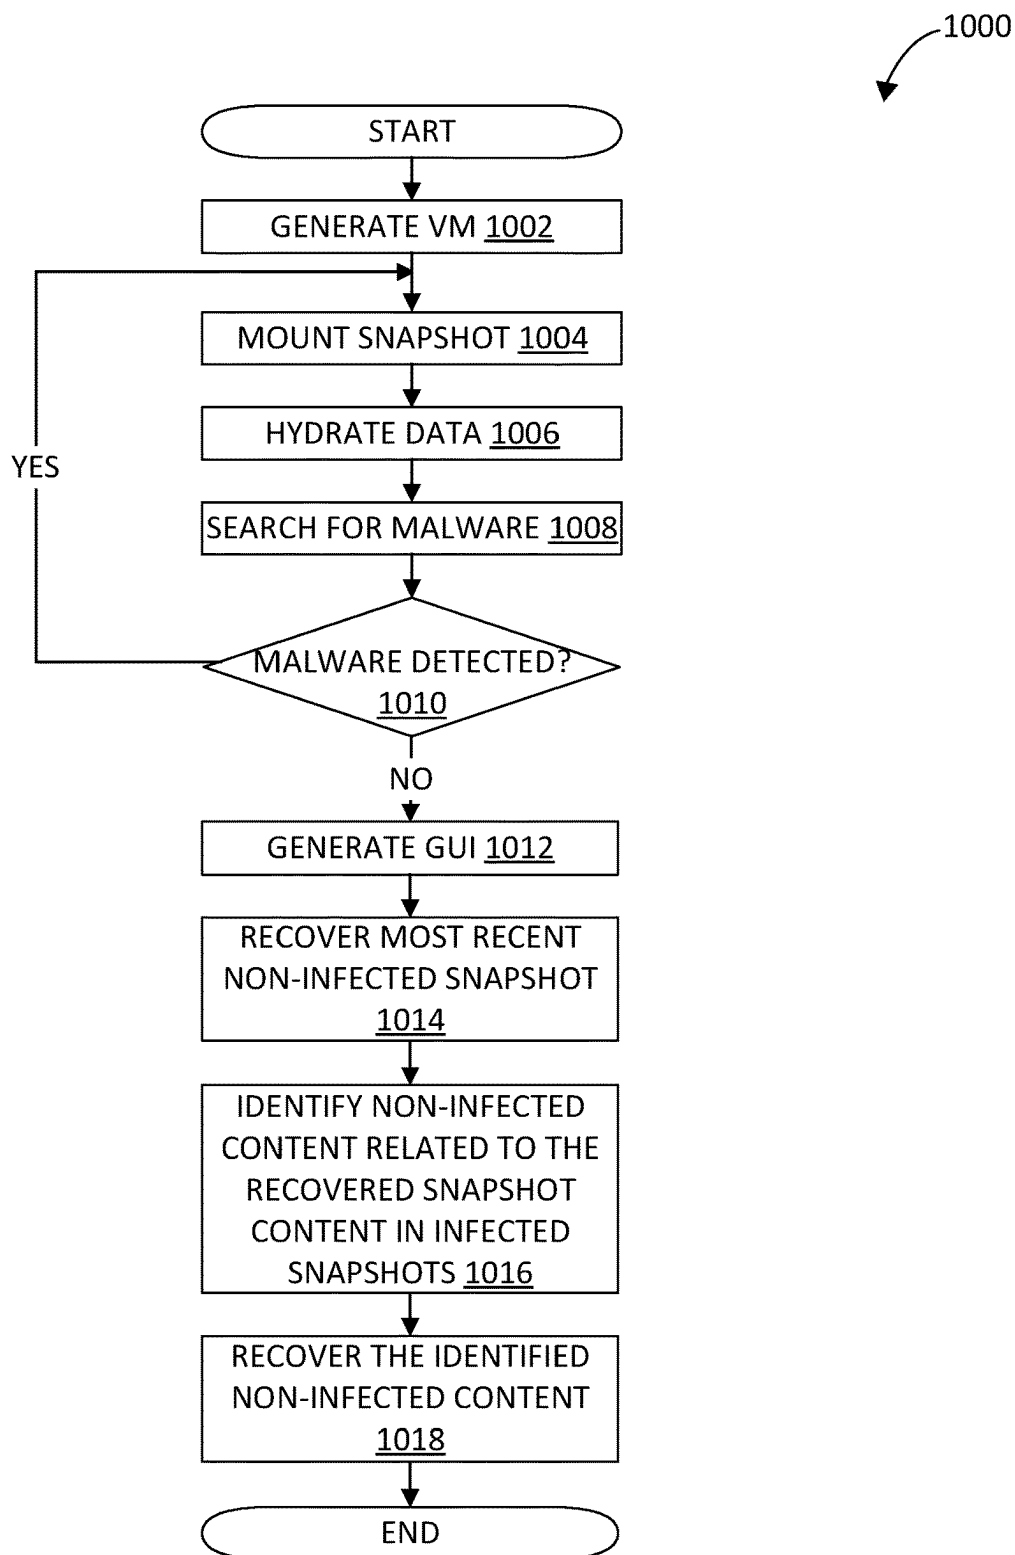


FIG. 10

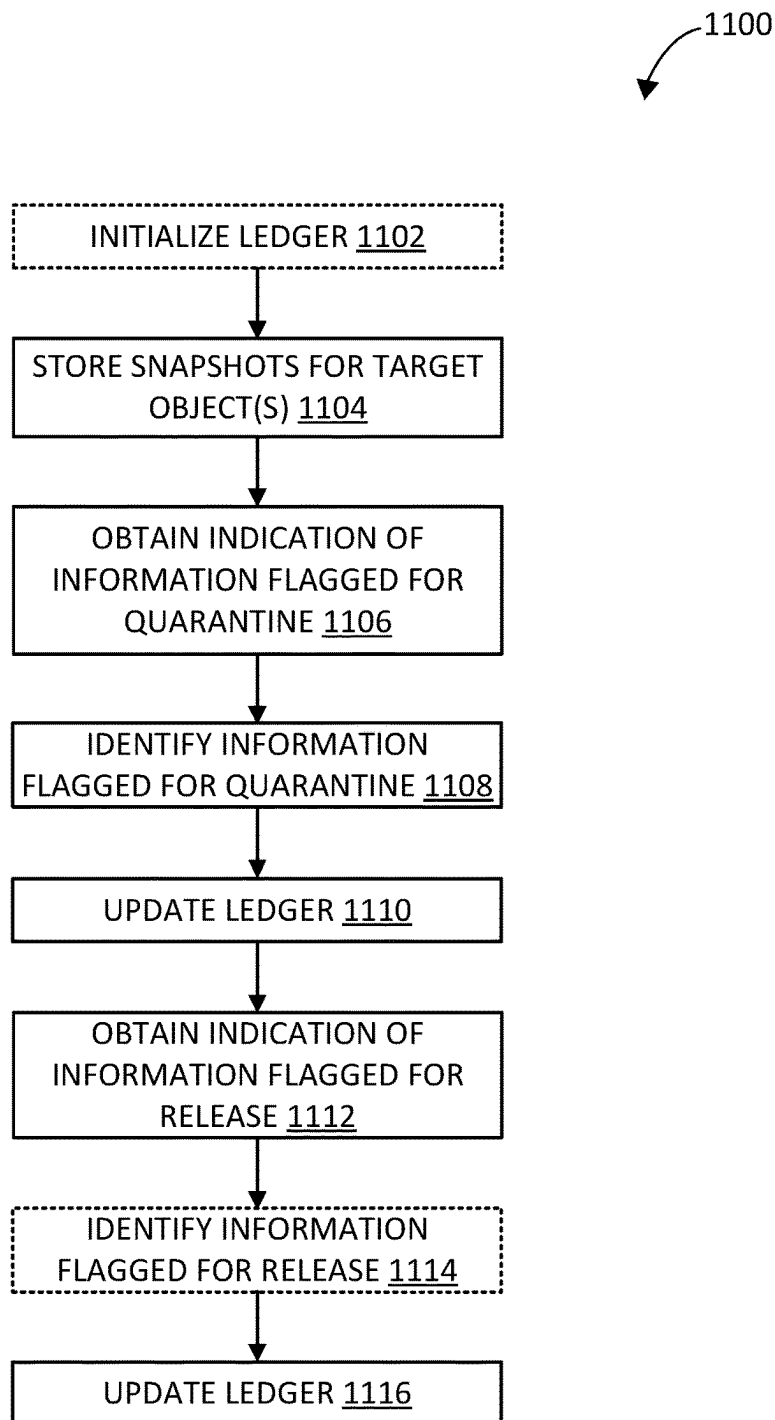


FIG. 11

1200

1205 Index	1210 Time	1215 Snapshot Identifier	1220 Path Identifier	1225 Operation Identifier
1	1/1/2021 9:00:00 AM	A	volume/usr/folder/file.ext	Quarantined
2	1/1/2021 9:01:00 AM	C	volume/usr/folder/*	Quarantined
3	1/1/2021 9:01:30 AM	E	volume/usr/folder/file.ext	Quarantined
4	1/1/2021 9:02:00 AM	D	volume/usr/folder/file.ext	Quarantined
5	1/1/2021 9:02:30 AM	C	volume/usr/folder2/file2.ext	Quarantined
6	1/1/2021 9:01:30 AM	E	volume/usr/folder/file2.ext	Quarantined
⋮				
10	1/2/2021 10:00:00 AM	A	volume/usr/folder/file.ext	Released
11	1/2/2021 10:01:00 AM	D	volume/usr/folder/file.ext	Released
12	1/2/2021 10:30:00 AM	C	volume/usr/folder/*	Released
13	1/2/2021 10:31:00 AM	E	volume/*	Released
⋮				
N	1/5/2021 8:00:00 AM	F	volume/usr/folder/file.ext	Quarantined
⋮				
⋮				

FIG. 12

1300

1315	1310	1320	1325
Snapshot Identifier	Time	Path Identifier	Quarantine Status
A	1/1/2021 9:00:00 AM	volume/usr/folder/file.ext	Quarantined
B	1/1/2021 9:01:00 AM	volume/usr/folder/*	Released
C	1/1/2021 9:01:00 AM	volume/usr/folder/*	Accessible
D	1/1/2021 9:01:00 AM	volume/usr/folder/*	Accessible
E	1/1/2021 9:01:30 AM	volume/usr/folder/file.ext	Quarantined
⋮			
H	1/5/2021 8:00:00 AM	volume/usr/folder/file.ext	Quarantined
⋮			

FIG. 13

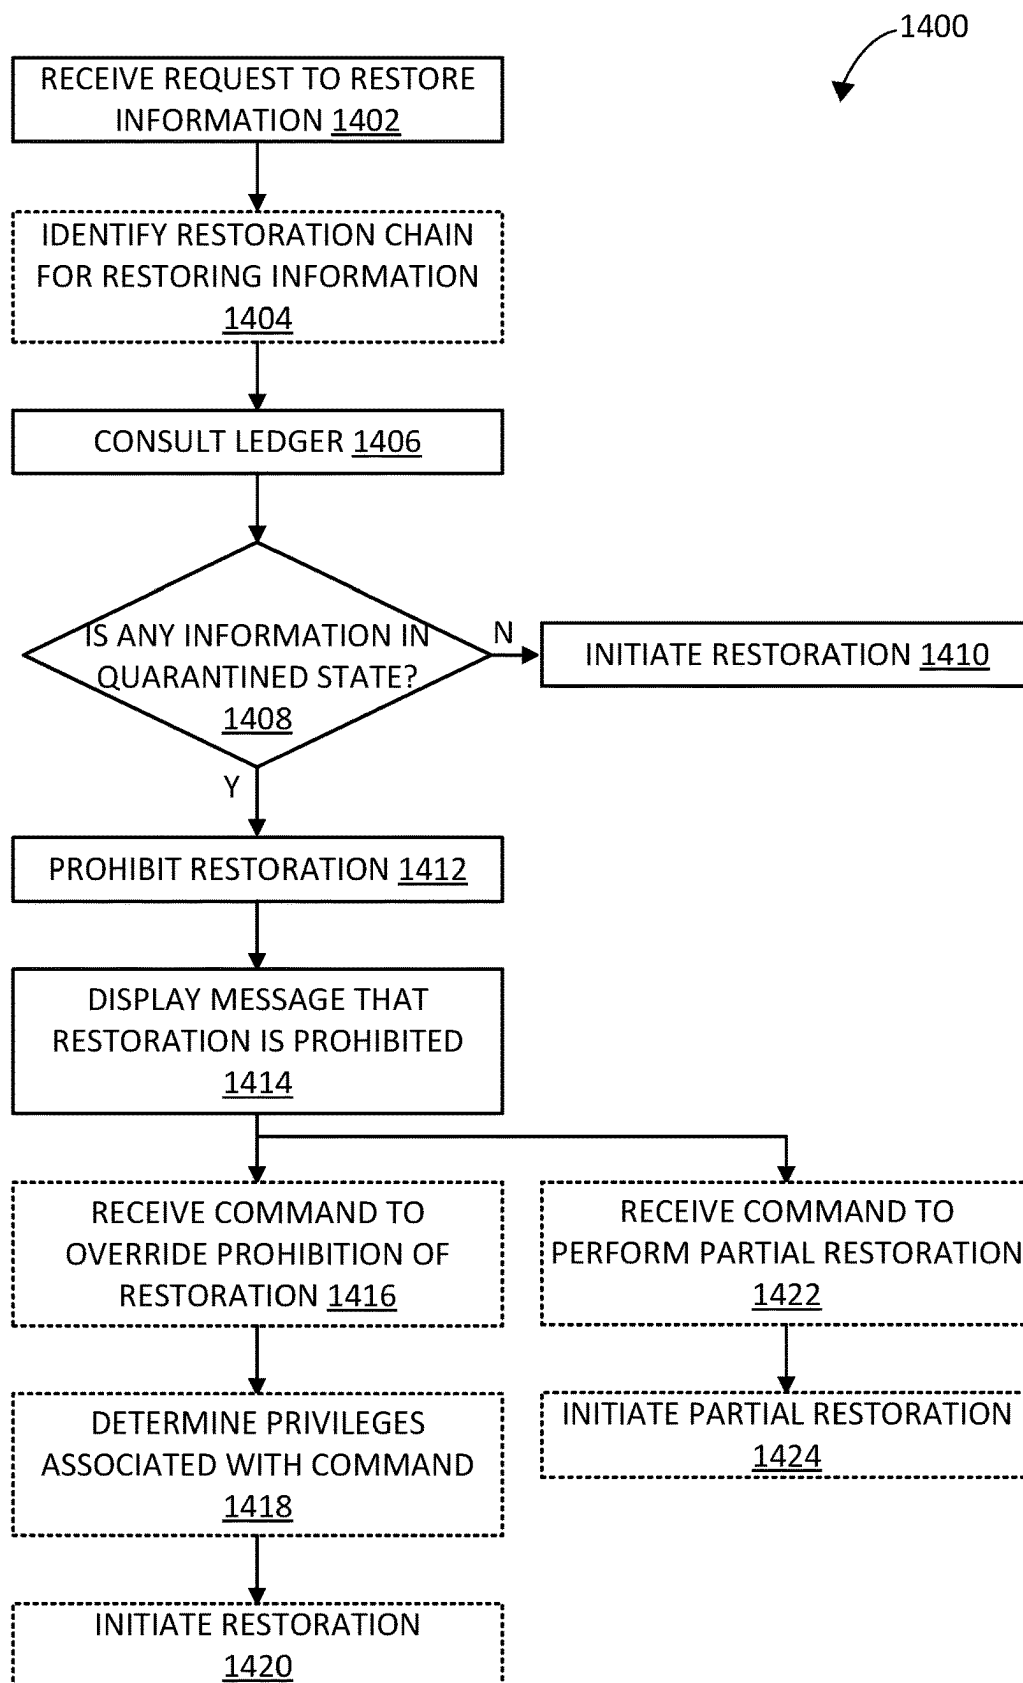


FIG. 14

1500

1505 1510 1515

Target Identifier	Snapshot Identifier	Quarantine Status	Path Identifier	...
1	A	Not Quarantined		
1	B	Not Quarantined		
1	C	Not Quarantined		
1	D	Quarantined	volume/usr/folder/file.ext	

⋮

1	J	Not Quarantined		
2	A	Not Quarantined		
2	B	Not Quarantined		
2	C	Not Quarantined		
3	A	Not Quarantined		
3	B	Not Quarantined		
3	C	Not Quarantined		
4	A	Not Quarantined		
4	B	Quarantined	volume/usr/folder/file.ext	
4	C	Not Quarantined		

⋮

N	A	Not Quarantined		
N	B	Quarantined	volume/usr/folder	
N	C	Released	volume/usr	

FIG. 15

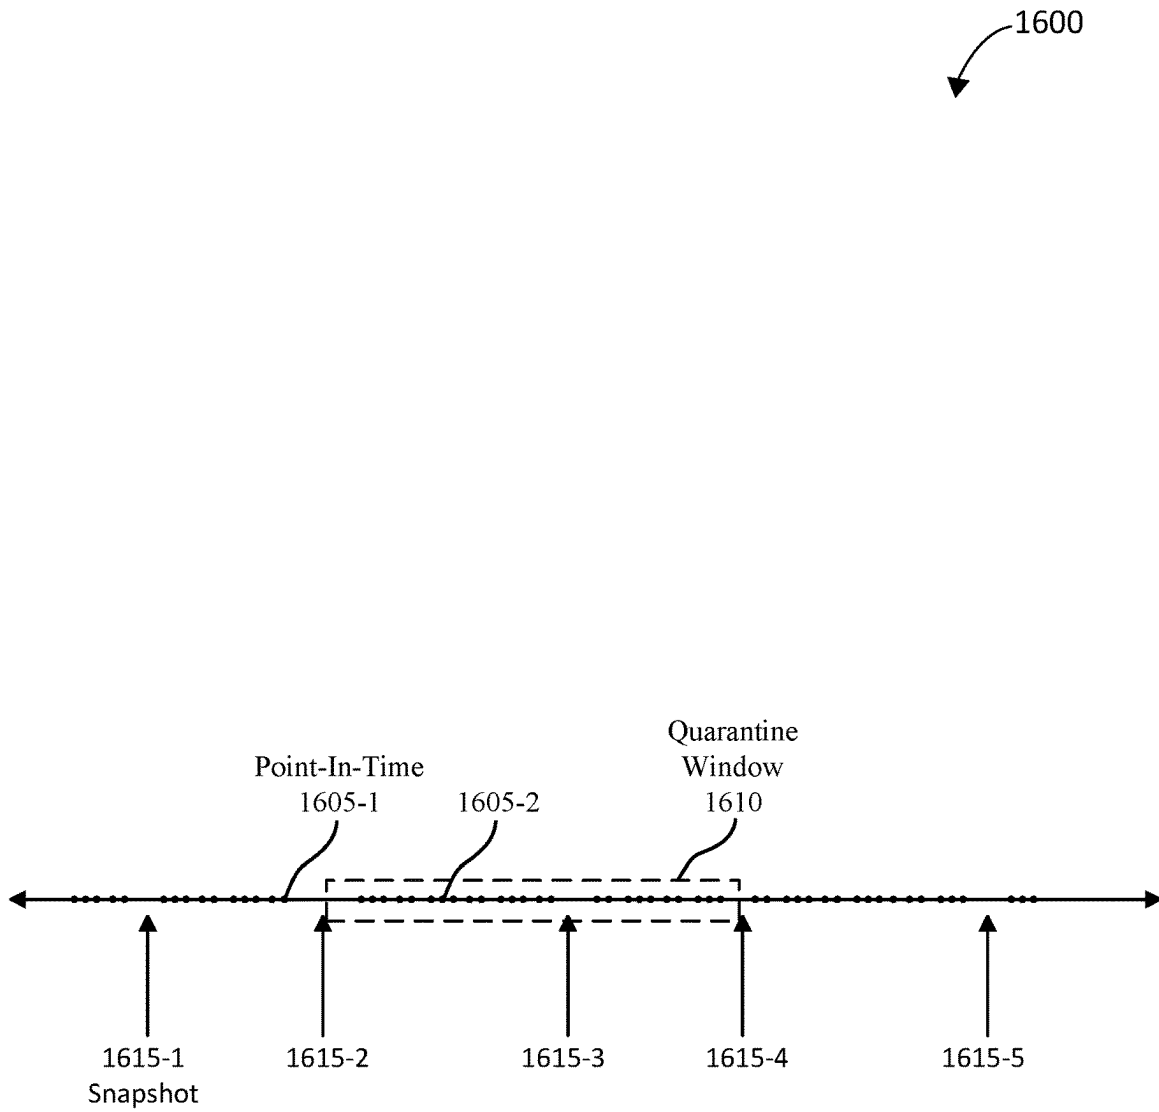


FIG. 16

1

RECOVERING QUARANTINED INFORMATION FROM BACKUP LOCATIONS

CROSS REFERENCE

The Present Application for Patent claims the benefit of U.S. Provisional Application No. 63/421,536 by Chandra et al., entitled "BULK SNAPSHOT RECOVERY" and filed Nov. 1, 2022; U.S. Provisional Application No. 63/319,953 by Chandra et al., entitled "QUARANTINING INFORMATION IN BACKUP LOCATIONS" and filed Mar. 15, 2022; and U.S. Provisional Application No. 63/276,822 by Gee et al., entitled "MALWARE DETECTION IN SNAPSHOTS" and filed Nov. 8, 2021, each of which is assigned to the assignee hereof and expressly incorporated by reference herein.

TECHNICAL FIELD

The present disclosure relates generally to data management including techniques for recovering quarantined information from backup locations.

BACKGROUND

The volume and complexity of data that is collected, analyzed, and stored is increasing rapidly over time. The computer infrastructure used to handle this data is also becoming more complex, with more processing power and more portability. As a result, data management and storage is becoming increasingly important. Significant issues of these processes include access to reliable data backup and storage, and fast data recovery in cases of failure. Other aspects include data portability across locations and plat-

BRIEF DESCRIPTION OF THE DRAWINGS

FIG. 1 depicts one embodiment of a networked computing environment in which the disclosed technology may be practiced, according to an example embodiment.

FIG. 2 depicts one embodiment of the server of FIG. 1, according to an example embodiment.

FIG. 3 depicts one embodiment of the storage appliance of FIG. 1, according to an example embodiment.

FIG. 4 shows an example cluster of a distributed decentralized database, according to some example embodiments.

FIG. 5 depicts a block diagram of a malware engine according to an example embodiment.

FIG. 6 depicts a flow chart illustrating a method of scanning a snapshot for malware according to an example embodiment.

FIG. 7 depicts an example interface according to an example embodiment.

FIG. 8 depicts a flow chart illustrating a method of recovering a non-infected file in an infected snapshot according to an example embodiment.

FIG. 9 depicts a flow chart illustrating a method of recovering an infected snapshot according to an example embodiment.

FIG. 10 depicts a flow chart illustrating a method of recovering non-infected content within an infected snapshot according to an example embodiment.

FIG. 11 depicts a flow chart illustrating a method of quarantining information in a snapshot according to an example embodiment.

2

FIGS. 12 and 13 depict example ledgers that support recording quarantining and release operations according to an example embodiment.

FIG. 14 depicts a flow chart illustrating a method of restoring a requested snapshot in accordance with quarantine information according to an example embodiment.

FIG. 15 depicts an example ledger that supports recording quarantining and release operations according to an example embodiment.

FIG. 16 depicts an example recovery timeline that supports recording quarantining and release operations according to an example embodiment.

DETAILED DESCRIPTION

The description that follows includes systems, methods, techniques, instruction sequences, and computing machine program products that embody illustrative embodiments of the present disclosure. In the following description, for purposes of explanation, numerous specific details are set forth in order to provide a thorough understanding of example embodiments. It will be evident, however, to one skilled in the art that the present inventive subject matter may be practiced without these specific details.

A portion of the disclosure of this patent document contains material that is subject to copyright protection. The copyright owner has no objection to the facsimile reproduction by anyone of the patent document or the patent disclosure, as it appears in the Patent and Trademark Office patent files or records, but otherwise reserves all copyright rights whatsoever. The following notice applies to the software and data as described below and in the drawings that form a part of this document: Copyright Rubrik, Inc., 2018-2021, All Rights Reserved.

It will be appreciated that some of the examples disclosed herein are described in the context of virtual machines that are backed up by using base and incremental snapshots, for example. This should not necessarily be regarded as limiting of the disclosures. The disclosures, systems and methods described herein apply not only to virtual machines of all types that run a file system (for example), but also to NAS devices, physical machines (for example Linux servers), and databases.

FIG. 1 depicts one embodiment of a networked computing environment **100** in which the disclosed technology may be practiced. As depicted, the networked computing environment **100** includes a data center **106**, a storage appliance **102**, and a computing device **108** in communication with each other via one or more networks **128**. The networked computing environment **100** may also include a plurality of computing devices interconnected through one or more networks **128**. The one or more networks **128** may allow computing devices and/or storage devices to connect to and communicate with other computing devices and/or other storage devices. In some cases, the networked computing environment **100** may include other computing devices and/or other storage devices not shown. The other computing devices may include, for example, a mobile computing device, a non-mobile computing device, a server, a workstation, a laptop computer, a tablet computer, a desktop computer, or an information processing system. The other storage devices may include, for example, a storage area network storage device, a networked-attached storage device, a hard disk drive, a solid-state drive, or a data storage system.

The data center **106** may include one or more servers, such as server **200**, in communication with one or more

3

storage devices, such as storage device **104**. The one or more servers may also be in communication with one or more storage appliances, such as storage appliance **102**. The server **200**, storage device **104**, and storage appliance **300** may be in communication with each other via a networking fabric connecting servers and data storage units within the data center **106** to each other. The storage appliance **300** may include a data management system for backing up virtual machines and/or files within a virtualized infrastructure. The server **200** may be used to create and manage one or more virtual machines associated with a virtualized infrastructure.

The one or more virtual machines may run various applications, such as a database application or a web server. The storage device **104** may include one or more hardware storage devices for storing data, such as a hard disk drive (HDD), a magnetic tape drive, a solid-state drive (SSD), a storage area network (SAN) storage device, or a Networked-Attached Storage (NAS) device. In some cases, a data center, such as data center **106**, may include thousands of servers and/or data storage devices in communication with each other. The one or more data storage devices **104** may comprise a tiered data storage infrastructure (or a portion of a tiered data storage infrastructure). The tiered data storage infrastructure may allow for the movement of data across different tiers of a data storage infrastructure between higher-cost, higher-performance storage devices (e.g., solid-state drives and hard disk drives) and relatively lower-cost, lower-performance storage devices (e.g., magnetic tape drives).

The one or more networks **128** may include a secure network such as an enterprise private network, an unsecure network such as a wireless open network, a local area network (LAN), a wide area network (WAN), and the Internet. The one or more networks **128** may include a cellular network, a mobile network, a wireless network, or a wired network. Each network of the one or more networks **128** may include hubs, bridges, routers, switches, and wired transmission media such as a direct-wired connection. The one or more networks **128** may include an extranet or other private network for securely sharing information or providing controlled access to applications or files.

A server, such as server **200**, may allow a client to download information or files (e.g., executable, text, application, audio, image, or video files) from the server **200** or to perform a search query related to particular information stored on the server **200**. In some cases, a server may act as an application server or a file server. In general, server **200** may refer to a hardware device that acts as the host in a client-server relationship or a software process that shares a resource with or performs work for one or more clients.

One embodiment of server **200** includes a network interface **110**, processor **112**, memory **114**, disk **116**, and virtualization manager **118** all in communication with each other. Network interface **110** allows server **200** to connect to one or more networks **128**. Network interface **110** may include a wireless network interface and/or a wired network interface. Processor **112** allows server **200** to execute computer-readable instructions stored in memory **114** in order to perform processes described herein. Processor **112** may include one or more processing units, such as one or more CPUs and/or one or more GPUs. Memory **114** may comprise one or more types of memory (e.g., RAM, SRAM, DRAM, ROM, EEPROM, Flash, etc.). Disk **116** may include a hard disk drive and/or a solid-state drive. Memory **114** and disk **116** may comprise hardware storage devices.

The virtualization manager **118** may manage a virtualized infrastructure and perform management operations associ-

4

ated with the virtualized infrastructure. The virtualization manager **118** may manage the provisioning of virtual machines running within the virtualized infrastructure and provide an interface to computing devices interacting with the virtualized infrastructure. In one example, the virtualization manager **118** may set a virtual machine having a virtual disk into a frozen state in response to a snapshot request made via an application programming interface (API) by a storage appliance, such as storage appliance **300**. Setting the virtual machine into a frozen state may allow a point in time snapshot of the virtual machine to be stored or transferred. In one example, updates made to a virtual machine that has been set into a frozen state may be written to a separate file (e.g., an update file) while the virtual disk may be set into a read-only state to prevent modifications to the virtual disk file while the virtual machine is in the frozen state.

The virtualization manager **118** may then transfer data associated with the virtual machine (e.g., an image of the virtual machine or a portion of the image of the virtual disk file associated with the state of the virtual disk at the point in time it is frozen) to a storage appliance (for example, a storage appliance **102** or storage appliance **300** of FIG. 1, described further below) in response to a request made by the storage appliance. After the data associated with the point in time snapshot of the virtual machine has been transferred to the storage appliance **300** (for example), the virtual machine may be released from the frozen state (i.e., unfrozen) and the updates made to the virtual machine and stored in the separate file may be merged into the virtual disk file. The virtualization manager **118** may perform various virtual machine-related tasks, such as cloning virtual machines, creating new virtual machines, monitoring the state of virtual machines, moving virtual machines between physical hosts for load balancing purposes, and facilitating backups of virtual machines.

One embodiment of a storage appliance **300** (or storage appliance **102**) includes a network interface **120**, processor **122**, memory **124**, and disk **126** all in communication with each other. Network interface **120** allows storage appliance **300** to connect to one or more networks **128**. Network interface **120** may include a wireless network interface and/or a wired network interface. Processor **122** allows storage appliance **300** to execute computer readable instructions stored in memory **124** in order to perform processes described herein. Processor **122** may include one or more processing units, such as one or more CPUs and/or one or more GPUs. Memory **124** may comprise one or more types of memory (e.g., RAM, SRAM, DRAM, ROM, EEPROM, NOR Flash, NAND Flash, etc.). Disk **126** may include a hard disk drive and/or a solid-state drive. Memory **124** and disk **126** may comprise hardware storage devices.

In one embodiment, the storage appliance **300** may include four machines. Each of the four machines may include a multi-core CPU, 64 GB of RAM, a 400 GB SSD, three 4 TB HDDs, and a network interface controller. In this case, the four machines may be in communication with the one or more networks **128** via the four network interface controllers. The four machines may comprise four nodes of a server cluster. The server cluster may comprise a set of physical machines that are connected together via a network. The server cluster may be used for storing data associated with a plurality of virtual machines, such as backup data associated with different point-in-time versions of the virtual machines.

The networked computing environment **100** may provide a cloud computing environment for one or more computing

5

devices. Cloud computing may refer to Internet-based computing, wherein shared resources, software, and/or information may be provided to one or more computing devices on-demand via the Internet. The networked computing environment **100** may comprise a cloud computing environment providing Software-as-a-Service (SaaS) or Infrastructure-as-a-Service (IaaS) services. SaaS may refer to a software distribution model in which applications are hosted by a service provider and made available to end users over the Internet. In one embodiment, the networked computing environment **100** may include a virtualized infrastructure that provides software, data processing, and/or data storage services to end users accessing the services via the networked computing environment **100**. In one example, networked computing environment **100** may provide cloud-based work productivity or business-related applications to a computing device, such as computing device **108**. The storage appliance **102** may comprise a cloud-based data management system for backing up virtual machines and/or files within a virtualized infrastructure, such as virtual machines running on server **200**/or files stored on server **200**.

In some cases, networked computing environment **100** may provide remote access to secure applications and files stored within data center **106** from a remote computing device, such as computing device **108**. The data center **106** may use an access control application to manage remote access to protected resources, such as protected applications, databases, or files located within the data center **106**. To facilitate remote access to secure applications and files, a secure network connection may be established using a virtual private network (VPN). A VPN connection may allow a remote computing device, such as computing device **108**, to securely access data from a private network (e.g., from a company file server or mail server) using an insecure public network or the Internet. The VPN connection may require client-side software (e.g., running on the remote computing device) to establish and maintain the VPN connection. The VPN client software may provide data encryption and encapsulation prior to the transmission of secure private network traffic through the Internet.

In some embodiments, the storage appliance **300** may manage the extraction and storage of virtual machine snapshots associated with different point in time versions of one or more virtual machines running within the data center **106**. A snapshot of a virtual machine may correspond with a state of the virtual machine at a particular point-in-time. In response to a restore command from the storage device **104**, the storage appliance **300** may restore a point-in-time version of a virtual machine or restore point-in-time versions of one or more files located on the virtual machine and transmit the restored data to the server **200**. In response to a mount command from the server **200**, the storage appliance **300** may allow a point-in-time version of a virtual machine to be mounted and allow the server **200** to read and/or modify data associated with the point-in-time version of the virtual machine. To improve storage density, the storage appliance **300** may deduplicate and compress data associated with different versions of a virtual machine and/or deduplicate and compress data associated with different virtual machines. To improve system performance, the storage appliance **300** may first store virtual machine snapshots received from a virtualized environment in a cache, such as a flash-based cache. The cache may also store popular data or frequently accessed data (e.g., based on a history of virtual machine restorations, incremental files associated with commonly restored virtual machine versions) and cur-

6

rent day incremental files or incremental files corresponding with snapshots captured within the past 24 hours.

An incremental file may comprise a forward incremental file or a reverse incremental file. A forward incremental file may include a set of data representing changes that have occurred since an earlier point-in-time snapshot of a virtual machine. To generate a snapshot of the virtual machine corresponding with a forward incremental file, the forward incremental file may be combined with an earlier point in time snapshot of the virtual machine (e.g., the forward incremental file may be combined with the last full image of the virtual machine that was captured before the forward incremental file was captured and any other forward incremental files that were captured subsequent to the last full image and prior to the forward incremental file). A reverse incremental file may include a set of data representing changes from a later point-in-time snapshot of a virtual machine. To generate a snapshot of the virtual machine corresponding with a reverse incremental file, the reverse incremental file may be combined with a later point-in-time snapshot of the virtual machine (e.g., the reverse incremental file may be combined with the most recent snapshot of the virtual machine and any other reverse incremental files that were captured prior to the most recent snapshot and subsequent to the reverse incremental file).

The storage appliance **300** may provide a user interface (e.g., a web-based interface or a graphical user interface) that displays virtual machine backup information such as identifications of the virtual machines protected and the historical versions or time machine views for each of the virtual machines protected. A time machine view of a virtual machine may include snapshots of the virtual machine over a plurality of points in time. Each snapshot may comprise the state of the virtual machine at a particular point in time. Each snapshot may correspond with a different version of the virtual machine (e.g., Version 1 of a virtual machine may correspond with the state of the virtual machine at a first point in time and Version 2 of the virtual machine may correspond with the state of the virtual machine at a second point in time subsequent to the first point in time).

The user interface may enable an end user of the storage appliance **300** (e.g., a system administrator or a virtualization administrator) to select a particular version of a virtual machine to be restored or mounted. When a particular version of a virtual machine has been mounted, the particular version may be accessed by a client (e.g., a virtual machine, a physical machine, or a computing device) as if the particular version was local to the client. A mounted version of a virtual machine may correspond with a mount point directory (e.g., /snapshots/VM5Nersion23). In one example, the storage appliance **300** may run an NFS server and make the particular version (or a copy of the particular version) of the virtual machine accessible for reading and/or writing. The end user of the storage appliance **300** may then select the particular version to be mounted and run an application (e.g., a data analytics application) using the mounted version of the virtual machine. In another example, the particular version may be mounted as an iSCSI target.

FIG. 2 depicts one embodiment of server **200** of FIG. 1. The server **200** may comprise one server out of a plurality of servers that are networked together within a data center (e.g., data center **106**). In one example, the plurality of servers may be positioned within one or more server racks within the data center. As depicted, the server **200** includes hardware-level components and software-level components. The hardware-level components include one or more processors **202**, one or more memory **204**, and one or more

disks **206**. The software-level components include a hypervisor **208**, a virtualized infrastructure manager **222**, and one or more virtual machines, such as virtual machine **220**. The hypervisor **208** may comprise a native hypervisor or a hosted hypervisor. The hypervisor **208** may provide a virtual operating platform for running one or more virtual machines, such as virtual machine **220**. Virtual machine **220** includes a plurality of virtual hardware devices including a virtual processor **210**, a virtual memory **212**, and a virtual disk **214**. The virtual disk **214** may comprise a file stored within the one or more disks **206**. In one example, a virtual machine **220** may include a plurality of virtual disks **214**, with each virtual disk of the plurality of virtual disks **214** associated with a different file stored on the one or more disks **206**. Virtual machine **220** may include a guest operating system **216** that runs one or more applications, such as application **218**.

The virtualized infrastructure manager **222**, which may correspond with the virtualization manager **118** in FIG. **1**, may run on a virtual machine or natively on the server **200**. The virtual machine may, for example, be or include the virtual machine **220** or a virtual machine separate from the server **200**. Other arrangements are possible. The virtualized infrastructure manager **222** may provide a centralized platform for managing a virtualized infrastructure that includes a plurality of virtual machines. The virtualized infrastructure manager **222** may manage the provisioning of virtual machines running within the virtualized infrastructure and provide an interface to computing devices interacting with the virtualized infrastructure. The virtualized infrastructure manager **222** may perform various virtualized infrastructure related tasks, such as cloning virtual machines, creating new virtual machines, monitoring the state of virtual machines, and facilitating backups of virtual machines.

In one embodiment, the server **200** may use the virtualized infrastructure manager **222** to facilitate backups for a plurality of virtual machines (e.g., eight different virtual machines) running on the server **200**. Each virtual machine running on the server **200** may run its own guest operating system and its own set of applications. Each virtual machine running on the server **200** may store its own set of files using one or more virtual disks associated with the virtual machine (e.g., each virtual machine may include two virtual disks that are used for storing data associated with the virtual machine).

In one embodiment, a data management application running on a storage appliance, such as storage appliance **102** in FIG. **1** or storage appliance **300** in FIG. **1**, may request a snapshot of a virtual machine running on server **200**. The snapshot of the virtual machine may be stored as one or more files, with each file associated with a virtual disk of the virtual machine. A snapshot of a virtual machine may correspond with a state of the virtual machine at a particular point in time. The particular point in time may be associated with a time stamp. In one example, a first snapshot of a virtual machine may correspond with a first state of the virtual machine (including the state of applications and files stored on the virtual machine) at a first point in time and a second snapshot of the virtual machine may correspond with a second state of the virtual machine at a second point in time subsequent to the first point in time.

In response to a request for a snapshot of a virtual machine at a particular point in time, the virtualized infrastructure manager **222** may set the virtual machine into a frozen state or store a copy of the virtual machine at the particular point in time. The virtualized infrastructure manager **222** may then transfer data associated with the virtual

machine (e.g., an image of the virtual machine or a portion of the image of the virtual machine) to the storage appliance **300** or storage appliance **102**. The data associated with the virtual machine may include a set of files including a virtual disk file storing contents of a virtual disk of the virtual machine at the particular point in time and a virtual machine configuration file storing configuration settings for the virtual machine at the particular point in time. The contents of the virtual disk file may include the operating system used by the virtual machine, local applications stored on the virtual disk, and user files (e.g., images and word processing documents). In some cases, the virtualized infrastructure manager **222** may transfer a full image of the virtual machine to the storage appliance **102** or storage appliance **300** of FIG. **1** or a plurality of data blocks corresponding with the full image (e.g., to enable a full image-level backup of the virtual machine to be stored on the storage appliance). In other cases, the virtualized infrastructure manager **222** may transfer a portion of an image of the virtual machine associated with data that has changed since an earlier point in time prior to the particular point in time or since a last snapshot of the virtual machine was taken. In one example, the virtualized infrastructure manager **222** may transfer only data associated with virtual blocks stored on a virtual disk of the virtual machine that have changed since the last snapshot of the virtual machine was taken. In one embodiment, the data management application may specify a first point in time and a second point in time and the virtualized infrastructure manager **222** may output one or more virtual data blocks associated with the virtual machine that have been modified between the first point in time and the second point in time.

In some embodiments, the server **200** or the hypervisor **208** may communicate with a storage appliance, such as storage appliance **102** in FIG. **1** or storage appliance **300** in FIG. **1**, using a distributed file system protocol such as Network File System (NFS) Version 3, or Server Message Block (SMB) protocol. The distributed file system protocol may allow the server **200** or the hypervisor **208** to access, read, write, or modify files stored on the storage appliance as if the files were locally stored on the server **200**. The distributed file system protocol may allow the server **200** or the hypervisor **208** to mount a directory or a portion of a file system located within the storage appliance.

FIG. **3** depicts one embodiment of storage appliance **300** in FIG. **1**. The storage appliance may include a plurality of physical machines that may be grouped together and presented as a single computing system. Each physical machine of the plurality of physical machines may comprise a node in a cluster (e.g., a failover cluster). In one example, the storage appliance may be positioned within a server rack within a data center. As depicted, the storage appliance **300** includes hardware-level components and software-level components. The hardware-level components include one or more physical machines, such as physical machine **314** and physical machine **324**. The physical machine **314** includes a network interface **316**, processor **318**, memory **320**, and disk **322** all in communication with each other. Processor **318** allows physical machine **314** to execute computer readable instructions stored in memory **320** to perform processes described herein. Disk **322** may include a hard disk drive and/or a solid-state drive. The physical machine **324** includes a network interface **326**, processor **328**, memory **330**, and disk **332** all in communication with each other. Processor **328** allows physical machine **324** to execute computer readable instructions stored in memory **330** to perform processes described herein. Disk **332** may include

a hard disk drive and/or a solid-state drive. In some cases, disk **332** may include a flash-based SSD or a hybrid HDD/SSD drive. In one embodiment, the storage appliance **300** may include a plurality of physical machines arranged in a cluster (e.g., eight machines in a cluster). Each of the plurality of physical machines may include a plurality of multi-core CPUs, 108 GB of RAM, a 500 GB SSD, four 4 TB HDDs, and a network interface controller.

In some embodiments, the plurality of physical machines may be used to implement a cluster-based network file-server. The cluster-based network file server may neither require nor use a front-end load balancer. One issue with using a front-end load balancer to host the IP address for the cluster-based network file server and to forward requests to the nodes of the cluster-based network file server is that the front-end load balancer comprises a single point of failure for the cluster-based network file server. In some cases, the file system protocol used by a server, such as server **200** in FIG. 1, or a hypervisor, such as hypervisor **208** in FIG. 2, to communicate with the storage appliance **300** may not provide a failover mechanism (e.g., NFS Version 3). In the case that no failover mechanism is provided on the client side, the hypervisor may not be able to connect to a new node within a cluster in the event that the node connected to the hypervisor fails.

In some embodiments, each node in a cluster may be connected to each other via a network and may be associated with one or more IP addresses (e.g., two different IP addresses may be assigned to each node). In one example, each node in the cluster may be assigned a permanent IP address and a floating IP address and may be accessed using either the permanent IP address or the floating IP address. In this case, a hypervisor, such as hypervisor **208** in FIG. 2, may be configured with a first floating IP address associated with a first node in the cluster. The hypervisor may connect to the cluster using the first floating IP address. In one example, the hypervisor may communicate with the cluster using the NFS Version 3 protocol. Each node in the cluster may run a Virtual Router Redundancy Protocol (VRRP) daemon. A daemon may comprise a background process. Each VRRP daemon may include a list of all floating IP addresses available within the cluster. In the event that the first node associated with the first floating IP address fails, one of the VRRP daemons may automatically assume or pick up the first floating IP address if no other VRRP daemon has already assumed the first floating IP address. Therefore, if the first node in the cluster fails or otherwise goes down, then one of the remaining VRRP daemons running on the other nodes in the cluster may assume the first floating IP address that is used by the hypervisor for communicating with the cluster.

In order to determine which of the other nodes in the cluster will assume the first floating IP address, a VRRP priority may be established. In one example, given a number (N) of nodes in a cluster from node(0) to node(N-1), for a floating IP address (i), the VRRP priority of nodeG) may be G-i) modulo N. In another example, given a number (N) of nodes in a cluster from node(0) to node(N-1), for a floating IP address (i), the VRRP priority of nodeG) may be (i-j) modulo N. In these cases, nodeG) will assume floating IP address (i) only if its VRRP priority is higher than that of any other node in the cluster that is alive and announcing itself on the network. Thus, if a node fails, then there may be a clear priority ordering for determining which other node in the cluster will take over the failed node's floating IP address.

In some cases, a cluster may include a plurality of nodes and each node of the plurality of nodes may be assigned a different floating IP address. In this case, a first hypervisor may be configured with a first floating IP address associated with a first node in the cluster, a second hypervisor may be configured with a second floating IP address associated with a second node in the cluster, and a third hypervisor may be configured with a third floating IP address associated with a third node in the cluster.

As depicted in FIG. 3, the software-level components of the storage appliance **300** may include data management system **302**, a virtualization interface **304**, a distributed job scheduler **308**, a distributed metadata store **310**, a distributed file system **312**, and one or more virtual machine search indexes, such as virtual machine search index **306**. In one embodiment, the software-level components of the storage appliance **300** may be run using a dedicated hardware-based appliance. In another embodiment, the software-level components of the storage appliance **300** may be run from the cloud (e.g., the software-level components may be installed on a cloud service provider).

In some cases, the data storage across a plurality of nodes in a cluster (e.g., the data storage available from the one or more physical machine (e.g., physical machine **314** and physical machine **324**)) may be aggregated and made available over a single file system namespace (e.g., /snapshots/). A directory for each virtual machine protected using the storage appliance **300** may be created (e.g., the directory for Virtual Machine A may be /snapshots/VM_A). Snapshots and other data associated with a virtual machine may reside within the directory for the virtual machine. In one example, snapshots of a virtual machine may be stored in subdirectories of the directory (e.g., a first snapshot of Virtual Machine A may reside in /snapshots/VM_A/s1/ and a second snapshot of Virtual Machine A may reside in /snapshots/VM_A/s2/).

The distributed file system **312** may present itself as a single file system, in which as new physical machines or nodes are added to the storage appliance **300**, the cluster may automatically discover the additional nodes and automatically increase the available capacity of the file system for storing files and other data. Each file stored in the distributed file system **312** may be partitioned into one or more chunks or shards. Each of the one or more chunks may be stored within the distributed file system **312** as a separate file. The files stored within the distributed file system **312** may be replicated or mirrored over a plurality of physical machines, thereby creating a load-balanced and fault tolerant distributed file system. In one example, storage appliance **300** may include ten physical machines arranged as a failover cluster and a first file corresponding with a snapshot of a virtual machine (e.g., /snapshots/VM_A/s1/s1.full) may be replicated and stored on three of the ten machines.

The distributed metadata store **310** may include a distributed database management system that provides high availability without a single point of failure. In one embodiment, the distributed metadata store **310** may comprise a database, such as a distributed document-oriented database. The distributed metadata store **310** may be used as a distributed key value storage system. In one example, the distributed metadata store **310** may comprise a distributed NoSQL key value store database. In some cases, the distributed metadata store **310** may include a partitioned row store, in which rows are organized into tables or other collections of related data held within a structured format within the key value store database. A table (or a set of tables) may be used to store metadata information associated with one or more files

11

stored within the distributed file system **312**. The metadata information may include the name of a file, a size of the file, file permissions associated with the file, when the file was last modified, and file mapping information associated with an identification of the location of the file stored within a cluster of physical machines.

In one embodiment, a new file corresponding with a snapshot of a virtual machine may be stored within the distributed file system **312** and metadata associated with the new file may be stored within the distributed metadata store **310**. The distributed metadata store **310** may also be used to store a backup schedule for the virtual machine and a list of snapshots for the virtual machine that are stored using the storage appliance **300**. In some examples, the metadata for a snapshot may include a time when the snapshot was taken, an expiration time for the snapshot, a quarantine status of the snapshot, and anomalous status of the snapshot (e.g., if malware is identified in the snapshot during ingestion, etc.).

In some cases, the distributed metadata store **310** may be used to manage one or more versions of a virtual machine. Each version of the virtual machine may correspond with a full image snapshot of the virtual machine stored within the distributed file system **312** or an incremental snapshot of the virtual machine (e.g., a forward incremental or reverse incremental) stored within the distributed file system **312**. In one embodiment, the one or more versions of the virtual machine may correspond with a plurality of files. The plurality of files may include a single full image snapshot of the virtual machine and one or more incremental aspects derived from the single full image snapshot. The single full image snapshot of the virtual machine may be stored using a first storage device of a first type (e.g., an HDD) and the one or more incremental aspects derived from the single full image snapshot may be stored using a second storage device of a second type (e.g., an SSD). In this case, only a single full image needs to be stored and each version of the virtual machine may be generated from the single full image or the single full image combined with a subset of the one or more incremental aspects. Furthermore, each version of the virtual machine may be generated by performing a sequential read from the first storage device (e.g., reading a single file from an HDD) to acquire the full image and, in parallel, performing one or more reads from the second storage device (e.g., performing fast random reads from an SSD) to acquire the one or more incremental aspects.

The distributed job scheduler **308** may be used for scheduling backup jobs that acquire and store virtual machine snapshots for one or more virtual machines over time. The distributed job scheduler **308** may follow a backup schedule to back up an entire image of a virtual machine at a particular point in time or one or more virtual disks associated with the virtual machine at the particular point in time. In one example, the backup schedule may specify that the virtual machine be backed up at a snapshot capture frequency, such as every two hours or every 24 hours. Each backup job may be associated with one or more tasks to be performed in a sequence. Each of the one or more tasks associated with a job may be run on a particular node within a cluster. In some cases, the distributed job scheduler **308** may schedule a specific job to be run on a particular node based on data stored on the particular node. For example, the distributed job scheduler **308** may schedule a virtual machine snapshot job to be run on a node in a cluster that is used to store snapshots of the virtual machine in order to reduce network congestion.

The distributed job scheduler **308** may comprise a distributed fault tolerant job scheduler, in which jobs affected

12

by node failures are recovered and rescheduled to be run on available nodes. In one embodiment, the distributed job scheduler **308** may be fully decentralized and implemented without the existence of a master node. The distributed job scheduler **308** may run job scheduling processes on each node in a cluster or on a plurality of nodes in the cluster. In one example, the distributed job scheduler **308** may run a first set of job scheduling processes on a first node in the cluster, a second set of job scheduling processes on a second node in the cluster, and a third set of job scheduling processes on a third node in the cluster. The first set of job scheduling processes, the second set of job scheduling processes, and the third set of job scheduling processes may store information regarding jobs, schedules, and the states of jobs using a metadata store, such as distributed metadata store **310**. In the event that the first node running the first set of job scheduling processes fails (e.g., due to a network failure or a physical machine failure), the states of the jobs managed by the first set of job scheduling processes may fail to be updated within a threshold period of time (e.g., a job may fail to be completed within 30 seconds or within minutes from being started). In response to detecting jobs that have failed to be updated within the threshold period of time, the distributed job scheduler **308** may undo and restart the failed jobs on available nodes within the cluster.

The job scheduling processes running on at least a plurality of nodes in a cluster (e.g., on each available node in the cluster) may manage the scheduling and execution of a plurality of jobs. The job scheduling processes may include run processes for running jobs, cleanup processes for cleaning up failed tasks, and rollback processes for rolling-back or undoing any actions or tasks performed by failed jobs. In one embodiment, the job scheduling processes may detect that a particular task for a particular job has failed and in response may perform a cleanup process to clean up or remove the effects of the particular task and then perform a rollback process that processes one or more completed tasks for the particular job in reverse order to undo the effects of the one or more completed tasks. Once the particular job with the failed task has been undone, the job scheduling processes may restart the particular job on an available node in the cluster.

The distributed job scheduler **308** may manage a job in which a series of tasks associated with the job are to be performed atomically (i.e., partial execution of the series of tasks is not permitted). If the series of tasks cannot be completely executed or there is any failure that occurs to one of the series of tasks during execution (e.g., a hard disk associated with a physical machine fails or a network connection to the physical machine fails), then the state of a data management system may be returned to a state as if none of the series of tasks was ever performed. The series of tasks may correspond with an ordering of tasks for the series of tasks and the distributed job scheduler **308** may ensure that each task of the series of tasks is executed based on the ordering of tasks. Tasks that do not have dependencies with each other may be executed in parallel.

In some cases, the distributed job scheduler **308** may schedule each task of a series of tasks to be performed on a specific node in a cluster. In other cases, the distributed job scheduler **308** may schedule a first task of the series of tasks to be performed on a first node in a cluster and a second task of the series of tasks to be performed on a second node in the cluster. In these cases, the first task may have to operate on a first set of data (e.g., a first file stored in a file system) stored on the first node and the second task may have to operate on a second set of data (e.g., metadata related to the

13

first file that is stored in a database) stored on the second node. In some embodiments, one or more tasks associated with a job may have an affinity to a specific node in a cluster.

In one example, if the one or more tasks require access to a database that has been replicated on three nodes in a cluster, then the one or more tasks may be executed on one of the three nodes. In another example, if the one or more tasks require access to multiple chunks of data associated with a virtual disk that has been replicated over four nodes in a cluster, then the one or more tasks may be executed on one of the four nodes. Thus, the distributed job scheduler **308** may assign one or more tasks associated with a job to be executed on a particular node in a cluster based on the location of data required to be accessed by the one or more tasks.

In one embodiment, the distributed job scheduler **308** may manage a first job associated with capturing and storing a snapshot of a virtual machine periodically (e.g., every 30 minutes). The first job may include one or more tasks, such as communicating with a virtualized infrastructure manager, such as the virtualized infrastructure manager **222** in FIG. 2, to create a frozen copy of the virtual machine and to transfer one or more chunks (or one or more files) associated with the frozen copy to a storage appliance, such as storage appliance **300** in FIG. 1. The one or more tasks may also include generating metadata for the one or more chunks, storing the metadata using the distributed metadata store **310**, storing the one or more chunks within the distributed file system **312**, and communicating with the virtualized infrastructure manager **222** that the frozen copy of the virtual machine may be unfrozen or released from a frozen state. The metadata for a first chunk of the one or more chunks may include information specifying a version of the virtual machine associated with the frozen copy, a time associated with the version (e.g., the snapshot of the virtual machine was taken at 5:30 p.m. on Jun. 29, 2018), and a file path to where the first chunk is stored within the distributed file system **92** (e.g., the first chunk is located at/snapshotsNM_B/s1/s1.chunk1). The one or more tasks may also include deduplication, compression (e.g., using a lossless data compression algorithm such as LZ4 or LZ77), decompression, encryption (e.g., using a symmetric key algorithm such as Triple DES or AES-256), and decryption related tasks.

The virtualization interface **304** may provide an interface for communicating with a virtualized infrastructure manager managing a virtualization infrastructure, such as virtualized infrastructure manager **222** in FIG. 2, and requesting data associated with virtual machine snapshots from the virtualization infrastructure. The virtualization interface **304** may communicate with the virtualized infrastructure manager using an Application Programming Interface (API) for accessing the virtualized infrastructure manager (e.g., to communicate a request for a snapshot of a virtual machine). In this case, storage appliance **300** may request and receive data from a virtualized infrastructure without requiring agent software to be installed or running on virtual machines within the virtualized infrastructure. The virtualization interface **304** may request data associated with virtual blocks stored on a virtual disk of the virtual machine that have changed since a last snapshot of the virtual machine was taken or since a specified prior point in time. Therefore, in some cases, if a snapshot of a virtual machine is the first snapshot taken of the virtual machine, then a full image of the virtual machine may be transferred to the storage appliance. However, if the snapshot of the virtual machine is not the first snapshot taken of the virtual machine, then only the

14

data blocks of the virtual machine that have changed since a prior snapshot was taken may be transferred to the storage appliance.

The virtual machine search index **306** may include a list of files that have been stored using a virtual machine and a version history for each of the files in the list. Each version of a file may be mapped to the earliest point-in-time snapshot of the virtual machine that includes the version of the file or to a snapshot of the virtual machine that includes the version of the file (e.g., the latest point in time snapshot of the virtual machine that includes the version of the file). In one example, the virtual machine search index **306** may be used to identify a version of the virtual machine that includes a particular version of a file (e.g., a particular version of a database, a spreadsheet, or a word processing document). In some cases, each of the virtual machines that are backed up or protected using storage appliance **300** may have a corresponding virtual machine search index.

In one embodiment, as each snapshot of a virtual machine is ingested, each virtual disk associated with the virtual machine is parsed in order to identify a file system type associated with the virtual disk and to extract metadata (e.g., file system metadata) for each file stored on the virtual disk. The metadata may include information for locating and retrieving each file from the virtual disk. The metadata may also include a name of a file, the size of the file, the last time at which the file was modified, and a content checksum for the file. Each file that has been added, deleted, or modified since a previous snapshot was captured may be determined using the metadata (e.g., by comparing the time at which a file was last modified with a time associated with the previous snapshot). Thus, for every file that has existed within any of the snapshots of the virtual machine, a virtual machine search index may be used to identify when the file was first created (e.g., corresponding with a first version of the file) and at what times the file was modified (e.g., corresponding with subsequent versions of the file). Each version of the file may be mapped to a particular version of the virtual machine that stores that version of the file.

In some cases, if a virtual machine includes a plurality of virtual disks, then a virtual machine search index may be generated for each virtual disk of the plurality of virtual disks. For example, a first virtual machine search index may catalog and map files located on a first virtual disk of the plurality of virtual disks and a second virtual machine search index may catalog and map files located on a second virtual disk of the plurality of virtual disks. In this case, a global file catalog or a global virtual machine search index for the virtual machine may include the first virtual machine search index and the second virtual machine search index. A global file catalog may be stored for each virtual machine backed up by a storage appliance within a file system, such as distributed file system **312** in FIG. 3.

The data management system **302** may comprise an application running on the storage appliance **300** that manages and stores one or more snapshots of a virtual machine. In one example, the data management system **302** may comprise a highest-level layer in an integrated software stack running on the storage appliance. The integrated software stack may include the data management system **302**, the virtualization interface **304**, the distributed job scheduler **308**, the distributed metadata store **310**, and the distributed file system **312**.

In some cases, the integrated software stack may run on other computing devices, such as a server or computing device **108** in FIG. 1. The data management system **302** may use the virtualization interface **304**, the distributed job

15

scheduler **308**, the distributed metadata store **310**, and the distributed file system **312** to manage and store one or more snapshots of a virtual machine. Each snapshot of the virtual machine may correspond with a point-in-time version of the virtual machine. The data management system **302** may generate and manage a list of versions for the virtual machine. Each version of the virtual machine may map to or reference one or more chunks and/or one or more files stored within the distributed file system **312**. Combined together, the one or more chunks and/or the one or more files stored within the distributed file system **312** may comprise a full image of the version of the virtual machine.

FIG. **4** shows an example cluster **400** of a distributed decentralized database, according to some example embodiments. As illustrated, the example cluster **400** includes five nodes, nodes 1-5. In some example embodiments, each of the five nodes runs from different machines, such as physical machine **314** in FIG. **3** or virtual machine **220** in FIG. **2**. The nodes in the example cluster **400** can include instances of peer nodes of a distributed database (e.g., cluster-based database, distributed decentralized database management system, a NoSQL database, Apache Cassandra, DataStax, MongoDB, CouchDB), according to some example embodiments. The distributed database system is distributed in that data is sharded or distributed across the example cluster **400** in shards or chunks and decentralized in that there is no central storage device and no single point of failure. The system operates under an assumption that multiple nodes may go down, up, or become non-responsive, and so on. Sharding is splitting up of the data horizontally and managing each shard separately on different nodes. For example, if the data managed by the example cluster **400** can be indexed using the 26 letters of the alphabet, node 1 can manage a first shard that handles records that start with A through E, node 2 can manage a second shard that handles records that start with F through J, and so on.

In some example embodiments, data written to one of the nodes is replicated to one or more other nodes per a replication protocol of the example cluster **400**. For example, data written to node 1 can be replicated to nodes 2 and 3. If node 1 prematurely terminates, node 2 and/or 3 can be used to provide the replicated data. In some example embodiments, each node of example cluster **400** frequently exchanges state information about itself and other nodes across the example cluster **400** using gossip protocol. Gossip protocol is a peer-to-peer communication protocol in which each node randomly shares (e.g., communicates, requests, transmits) location and state information about the other nodes in a given cluster.

Writing: For a given node, a sequentially written commit log captures the write activity to ensure data durability. The data is then written to an in-memory structure (e.g., a memtable, write-back cache). Each time the in-memory structure is full, the data is written to disk in a Sorted String Table data file. In some example embodiments, writes are automatically partitioned and replicated throughout the example cluster **400**.

Reading: Any node of example cluster **400** can receive a read request (e.g., query) from an external client. If the node that receives the read request manages the data requested, the node provides the requested data. If the node does not manage the data, the node determines which node manages the requested data. The node that received the read request then acts as a proxy between the requesting entity and the node that manages the data (e.g., the node that manages the data sends the data to the proxy node, which then provides the data to an external entity that generated the request).

16

The distributed decentralized database system is decentralized in that there is no single point of failure due to the nodes being symmetrical and seamlessly replaceable. For example, whereas conventional distributed data implementations have nodes with different functions (e.g., master/slave nodes, asymmetrical database nodes, federated databases), the nodes of example cluster **400** are configured to function the same way (e.g., as symmetrical peer database nodes that communicate via gossip protocol, such as Cassandra nodes) with no single point of failure. If one of the nodes in example cluster **400** terminates prematurely (“goes down”), another node can rapidly take the place of the terminated node without disrupting service. The example cluster **400** can be a container for a keyspace, which is a container for data in the distributed decentralized database system (e.g., whereas a database is a container for containers in conventional relational databases, the Cassandra keyspace is a container for a Cassandra database system).

FIG. **5** depicts a block diagram of a malware engine **502** according to an example embodiment. The malware engine **502** comprises a hydrator **504**, a mounter **506**, a detector **508**, a Yet Another Recursive/Ridiculous Acronym YARA rules **510**, hashes **512**, a flagger **514**, a recoverer **516**, a privilege determinator **518**, user privileges **520**, a propagator **522**, and a graphical user interface GUI **524**.

As will be discussed in more detail below, the malware engine **502** detects indicators of compromise that is present on a snapshot of an object (e.g., virtual machine, database, file system, etc.) that shows the snapshot may have been compromised by malware, such as ransomware. Ransomware is a piece of malware which infects an enterprise and encrypts its data. Embodiments enable an enterprise to quickly recover all protected objects to a safe copy, bringing the business back online as soon as possible, cutting out the malware from IT infrastructure, and restoring the maximum amount of data possible.

In order to initiate a recovery, the malware engine **502** determines the most recent point in time snapshot for each object that was not infected and enables push-button recovery of the determined snapshot using the GUI **524**. Specifically, a user can select an object or objects with the GUI **524** to scan snapshots in a snapshot chain of the object, e.g., starting with a most recent snapshot and scanning successively older snapshots (or skipping snapshots) in a snapshot chain (reverse chronological order or other order). Alternatively, a user can specify a range of snapshot to scan as well as specific objects, directories, etc. For each snapshot, the hydrator **504** will hydrate the snapshot (e.g., materialize/instantiate the snapshot e.g., via zero data copy) and the mounter **506** will mount (e.g., read without necessarily writing or transferring data as in a restore) the hydrated snapshot in a virtual machine (e.g., created by the hypervisor **208**), which may be sandboxed (e.g., no or limited network access) via, for example, user-mode Linux. The detector **508** then scans the mounted snapshot using YARA rules **510** and/or hashes **512** for malware. YARA rules **510** are a domain-specific language by which intelligence about indicators of compromise can be written and shared for threat hunting purposes. They typically allow for the specification of text or binary based indicators.

The detector **508** can be set to scan for malware based on all or a subset of the hashes **512** and/or the YARA rules **510**. For example, a query may be files with hashes A and B modified in the past month. Scanning may also be done by filename. After a snapshot is determined to be infected based on the presence of an indicator of compromise (e.g., matching hash and/or satisfied YARA rule), the GUI **524** can

17

display an interface showing infected versus non-infected snapshots as in example interface 700. The detector 508 can be deployed so that it examines all objects in a system and displays results for all objects as shown in the example interface 700.

The recoverer 516 recovers (e.g., restores, reads, mounts, etc.) non-infected snapshots, and, subject to user privileges 520, can recover infected snapshots and content (files) for forensic analysis, e.g., to a sandboxed virtual machine. The recoverer 516 can also restore non-infected files on infected snapshots. The privilege determinator 518 optionally limits what snapshots or content can be restored based on the user privileges 520. For example, by default, only non-infected snapshots may be restored. However, per user privileges 520, infected snapshots may be restored if a specific user requesting the restore has high enough privileges.

The propagator 522 propagates an infected (quarantine) status for a snapshot or content therein to other locations where a snapshot may be, e.g., archives or replicas of the snapshot, so that the duplicate of the infected snapshot isn't accidentally restored, thereby spreading malware.

FIG. 6 depicts a flowchart illustrating a method 600 of scanning a snapshot for malware according to an example embodiment. In an example embodiment, the storage appliance 300 can execute the method 600 using the malware engine 502. Example methods described herein may also be implemented in the form of executable instructions stored on a machine-readable medium or in the form of electronic circuitry. For instance, the operations of the method 600 may be represented by executable instructions that, when executed by a processor of a computing device, cause the computing device to perform the method 600. Depending on the embodiment, an operation of an example method described herein may be repeated in different ways or involve intervening operations not shown. Though the operations of example methods may be depicted and described in a certain order, the order in which the operations are performed may vary among embodiments, including performing certain operations in parallel.

At operation 602, the hypervisor 208 generates a virtual machine, e.g., virtual machine 220, which can be sandboxed (e.g., no or limited network access to prevent the spread of malware to other parts of a network). A snapshot is then mounted at operation 604, e.g., with Rubrik's LIVE-MOUNT. The snapshot may be part of a snapshot chain and accordingly, a most recent snapshot may be mounted or a specific snapshot selected for mounting. A range of snapshots can be specified as well as, snapshots of objects and/or directories. Data within the snapshot is then hydrated with the hydrator 504 at operation 606, which can include materializing/instantiating the snapshot.

The detector 508 then searches the mounted hydrated snapshot for malware at operation 608. Searching can include applying YARA rules 510 and/or looking for matching hashes 512 (e.g., MD5, SHA1, and/or SHA256) and/or looking for specific file names (e.g., by name, such as ABAP/Rivpas.c.intd, and/or by path prefix) and/or path matches. The YARA rules 510 and/or hashes 512 may be imported (e.g., if new YARA rules 510 or hashes 512 are created) and the detector 508 can search for malware 608 using all or a subset of the YARA rules 510 and/or hashes 512.

For example, given a set of Indicators of Compromise encoded as YARA rules in a CISA alert (such as Alert (AA20-302A) Ransomware Activity Targeting the Healthcare and Public Health Sector), the detector 508 identifies any snapshots where those indicators were found. In some

18

examples, the detector 508 identifies files associated with the indicators. Additionally, or alternatively, the detector 508 may receive a set of suspicious file hashes or file paths (such as from Microsoft's Hafnium IOC list) and may identify snapshots including the suspicious file hashes or file paths. In some examples, the detector 508 may analyze each file in a snapshot using one or more of: the YARA rules, suspicious file hashes, and suspicious file paths. Accordingly, the detector 508 may detect malware (e.g., infected files), maliciously encrypted data, or both. Additionally, or alternatively, the detector 508 may analyze sectors of a snapshot (including multiple files) or an entire snapshot using one or more of: the YARA rules, suspicious file hashes, and suspicious file paths.

Further, the detector 508 may restrict a search to specific snapshots or range of snapshots, a subset of files/directories and/or scan a replica of a snapshot instead an original snapshot. Restricting the search can include specific subsets of files or directories; file size limits, file creation/modification timestamps (e.g., before, after, or between timestamp(s)); file ownership; and files added/modified in a current snapshot (i.e., files that were not in prior snapshot or modified since that snapshot). Note that once a file in a snapshot is determined to be infected, it can be inferred that all later snapshots are infected (even if outside of a set range) and/or that specific file in the snapshot will be infected in later snapshots.

If, at operation 610, malware is detected, then a next most recent snapshot in the snapshot chain is mounted and the process is repeated until no malware is detected at operation 610. Further, in case of false negatives, additional snapshots in the chain can be scanned after no malware is detected in operation 610. Alternatively, the mounting 604 through the detecting 610 can start with an oldest or base snapshot and assuming no malware was detected in the base snapshot, repeating the process until malware is detected. Further, once malware is detected in a snapshot, the remainder of the snapshot does not need to be scanned and the next snapshot in the snapshot chain can be searched for malware.

Further, the mounting 604 through the detecting 610 can be repeated for all or some objects of a system including virtual machines, file systems, databases, network attached storage, etc. When malware is detected, metadata in the snapshot can be adjusted to indicate quarantine status. For example, quarantine-related metadata for the infected snapshots (e.g., quarantine status metadata, among other quarantine-related metadata) may be updated for the infected snapshot. Additionally, or alternatively, a ledger (e.g., an append-only ledger, a snapshot-level ledger, etc.) used to keep track of snapshots that have been quarantined, released from quarantine, or both, may be maintained. In some examples, the quarantine related-metadata for a snapshot is updated by updating related fields in a snapshot-level quarantine metadata table. Quarantining a snapshot is described in more detail herein, including with reference to FIG. 11.

At operation 612, the GUI 524 generates an interface, such as example interface 700, as will be discussed in further detail in conjunction with FIG. 7. For each object, the example interface 700 can show snapshots over time with a status (e.g., color coded) of each snapshot (infected, not infected, encrypted, infected and partially encrypted). The example interface 700 may also illustrate a cut point above which snapshots are quarantined due to infection (e.g., not recoverable, recovery not permitted, recoverable with sufficient user privileges, partially recoverable, etc.). The GUI 524 may display object(s) scanned (VM, Host/Share/fileset, etc.); snapshot(s) scanned; Date/time of scan; scan filter

criteria; Hash/Rules, etc. being scanned for (e.g., for YARA Rules: i. Rule name ii. Namespace iii. Tags iv. Hash of the rule, etc.); number of matched files; and number of matches; etc.

The GUI **524** may also provide detailed results after a search including: 1. The file name & path where a match was found 2. The rule (hash, or YARA name, namespace and tags) that were matched 3. The time when this file was created 4. The time when this file was modified 5. The owner of this file a. Their fully qualified name (if retrievable) b. Any security identifier (e.g., SID for Windows, User ID for Linux/Unix) 6. The MD5, SHA1, and SHA256 hashes of this file, etc.

A user then, using an interface such as the example interface **700**, enters a command to recover a snapshot, which is received at operation **614**. If the selected snapshot is determined to be quarantined at operation **616**, then the method **600** ends. Else, the snapshot can be recovered by the recoverer **516** at operation **618**, which can include mounting and/or restoring, etc. to a specified destination. Recovering snapshots in the context of quarantining is described in more detail herein, including with reference to FIG. 13.

Accordingly, the method **600** enables users to recover data in infected system without compromising recovered systems. For example, the method **600** enables users to restore an object to a point prior to a malware infection by quickly identifying healthy snapshots that can be used to perform a full system restore. Further, as the method **600** quarantines infected snapshots, the method **600** prevents reinfection by malware.

FIG. 7 depicts an example interface **700** according to an example embodiment. The GUI **524** generates the example interface **700** in one embodiment. The example interface **700** illustrates, for each object, a snapshot chain (e.g., in chronological order) and the status of each snapshot in each chain (or for the specified snapshots or range scanned). Further, the example interface **700** may illustrate a cut point indicating a quarantining of snapshots due to infection. For example, snapshots **702** and **708** are below the cut point and therefore not infected and can be restored and not quarantined. On the other hand, snapshot **704** is infected with malware while snapshot **706** is infected with malware and partially encrypted by that malware. On the other hand, snapshot **710** is fully encrypted by malware (which can be determined via entropy measurement).

FIG. 8 depicts a flowchart illustrating a method **800** of recovering a non-infected file in an infected snapshot according to an example embodiment. In an example embodiment, the storage appliance **300** can execute the method **800** using the malware engine **502**. Example methods described herein may also be implemented in the form of executable instructions stored on a machine-readable medium or in the form of electronic circuitry. For instance, the operations of the method **800** may be represented by executable instructions that, when executed by a processor of a computing device, cause the computing device to perform the method **800**. Depending on the embodiment, an operation of an example method described herein may be repeated in different ways or involve intervening operations not shown. Though the operations of example methods may be depicted and described in a certain order, the order in which the operations are performed may vary among embodiments, including performing certain operations in parallel.

At operation **802**, the hypervisor **208** generates a virtual machine, e.g., virtual machine **220**, which can be sandboxed (e.g., no or minimal network access to prevent the spread of

malware to other parts of a network). A snapshot is then mounted at operation **804**, e.g., with Rubrik's LIVE-MOUNT. The snapshot may be part of a snapshot chain and accordingly, a most recent snapshot may be mounted or a specific snapshot selected for mounting. Data within the snapshot is then hydrated with the hydrator **504** at operation **806**. The detector **508** then searches the mounted hydrated snapshot for malware at operation **808**. Searching can include applying YARA rules **510** and/or looking for matching hashes **512** (e.g., MD5, SHA1, and/or SHA256) and/or looking for specific file names (e.g., by name, such as ABAP/Rivpas.c.intd, and/or by path prefix) and/or path matches. The YARA rules **510** and/or hashes **512** may be imported (e.g., if new YARA rules **510** or hashes **512** are created) and the detector **508** can search for malware **808** using all or a subset of the YARA rules **510** and/or the hashes **512**.

For example, given a set of Indicators of Compromise encoded as YARA rules in a CISA alert (such as Alert (AA20-302A) Ransomware Activity Targeting the Healthcare and Public Health Sector), the detector **508** may identify any snapshots, file paths, or both, where those indicators were found. Or, given a set of suspicious file hashes or file paths (such as from Microsoft's Hafnium IOC list), the detector **508** may identify those snapshots, file paths, or both as well.

Further, the detector **508** may restrict a search to a range of snapshots, a subset of files/directories and/or scan a replica of a snapshot instead an original snapshot. Restricting the search can include specific snapshots or range of snapshots, specific subsets of files or directories; file size limits, file creation/modification timestamps (e.g., before, after between timestamp(s)); file ownership; and files added/modified in a current snapshot (i.e., files that were not in prior snapshot or modified since that snapshot). Note that once a file in a snapshot is determined to be infected, it can be inferred that all later snapshots are infected (even if outside of a set range) and/or that specific file in the snapshot will be infected in later snapshots.

If, at operation **810**, malware is detected, then a next most recent snapshot in the snapshot chain is mounted and the process is repeated until no malware is detected at operation **810**. Further, in case of false negatives, additional snapshots in the chain can be scanned after no malware is detected in operation **810**. Alternatively, the mounting **804** through the detecting **810** can start with an oldest or base snapshot and assuming no malware was detected in the base snapshot, repeating the process until malware is detected.

Further, the mounting **804** through the detecting **810** can be repeated for all objects of a system including virtual machines, file systems, databases, network attached storage, etc. When malware is detected, metadata in the snapshot can be adjusted to indicate quarantine status.

After the searching **808** is complete, the infected snapshots are quarantined at operation **812**, e.g., by adjusting metadata of the infected snapshots (e.g., changing a bit in the metadata for a field that indicates infected, encrypted, etc.). For example, quarantine-related metadata for the infected snapshots (e.g., quarantine status metadata, among other quarantine-related metadata) may be updated for the infected snapshot. Additionally, or alternatively, a ledger (e.g., an append-only ledger, a snapshot-level ledger, etc.) used to keep track of snapshots that have been quarantined, released from quarantine, or both, may be maintained. In some examples, the quarantine related-metadata for a snapshot is updated by updating related fields in a snapshot-level

21

quarantine metadata table. Quarantining a snapshot is described in more detail herein, including with reference to FIG. 11.

At operation 814, the GUI 524 generates an interface, such as the example interface 700. For each object, the example interface 700 can show snapshots over time with a status (e.g., color coded) of each snapshot (infected, not infected, encrypted, infected and partially encrypted). The example interface 700 may also illustrate a cut point above which snapshots are quarantined (e.g., not recoverable, recovery not permitted, recoverable with sufficient user privileges, partially recoverable, etc.). The GUI 524 may display object(s) scanned (VM, Host/Share/fileset, etc.); snapshot(s) scanned; Date/time of scan; scan filter criteria; Hash/Rules, etc. being scanned for (e.g., for YARA Rules: i. Rule name ii. Namespace iii. Tags iv. Hash of the rule, etc.); number of matched files; and number of matches; etc.

The GUI 524 may also provide detailed results after a search including: 1. The file name & path where a match was found 2. The rule (hash, or YARA name, namespace and tags) that were matched 3. The time when this file was created 4. The time when this file was modified 5. The owner of this file a. Their fully qualified name (if retrievable) b. Any security identifier (e.g., SID for Windows, User ID for Linux/Unix) 6. The MD5, SHA1, and SHA256 hashes of this file, etc.

A user then selects an infected snapshot, and the GUI 524 displays a list of non-infected files in the infected snapshot at operation 816 as determined by the search for malware at operation 808. Per a user command, a selected non-infected file from the infected snapshot is then recovered (e.g., mounted, restored, viewed, etc.) at operation 820 with the recoverer 516. The method 800 then ends.

FIG. 9 depicts a flowchart illustrating a method 900 of recovering an infected snapshot according to an example embodiment. In an example embodiment, the storage appliance 300 can execute the method 900 using the malware engine 502. Example methods described herein may also be implemented in the form of executable instructions stored on a machine-readable medium or in the form of electronic circuitry. For instance, the operations of the method 900 may be represented by executable instructions that, when executed by a processor of a computing device, cause the computing device to perform the method 900. Depending on the embodiment, an operation of an example method described herein may be repeated in different ways or involve intervening operations not shown. Though the operations of example methods may be depicted and described in a certain order, the order in which the operations are performed may vary among embodiments, including performing certain operations in parallel.

At operation 902, the hypervisor 208 generates a virtual machine, e.g., virtual machine 220, which can be sandboxed (e.g., no network access to prevent the spread of malware to other parts of a network). A snapshot is then mounted at operation 904, e.g., with Rubrik's LIVEMOUNT. The snapshot may be part of a snapshot chain and accordingly, a most recent snapshot may be mounted or a specific snapshot selected for mounting. Data within the snapshot is then hydrated with the hydrator 504 at operation 906. The detector 508 then searches the mounted hydrated snapshot for malware at operation 908. Searching can include applying YARA rules 510 and/or looking for matching hashes 512 (e.g., MD5, SHA1, and/or SHA256) and/or looking for specific file names (e.g., by name, such as ABAP/Rivpass.c.intd, and/or by path prefix) and/or path matches. The YARA rules 510 and/or hashes 512 may be imported (e.g.,

22

if new YARA rules 510 or hashes 512 are created) and the detector 508 can search for malware 908 using all or a subset of the YARA rules 510 and/or hashes 512.

For example, given a set of Indicators of Compromise encoded as YARA rules in a CISA alert (such as Alert (AA20-302A) Ransomware Activity Targeting the Healthcare and Public Health Sector), the detector 508 may identify any snapshots, file paths, or both, where those indicators were found. Or, given a set of suspicious file hashes or file paths (such as from Microsoft's Hafnium IOC list), the detector 508 may identify those snapshots, file paths, or both as well.

Further, the detector 508 may restrict a search to a specific snapshots or range of snapshots, a subset of files/directories and/or scan a replica of a snapshot instead an original snapshot. Restricting the search can include specific subsets of files or directories; file size limits, file creation/modification timestamps (e.g., before, after between timestamp(s)); file ownership; and files added/modified in a current snapshot (i.e., files that were not in prior snapshot or modified since that snapshot).

If, at operation 910, malware is detected, then a next most recent snapshot in the snapshot chain is mounted and the process is repeated until no malware is detected at operation 910. Further, in case of false negatives, additional snapshots in the chain can be scanned after no malware is detected in operation 910. Alternatively, the mounting 904 through the detecting 910 can start with an oldest or base snapshot and assuming no malware was detected in the base snapshot, repeating the process until malware is detected. Further, once malware is detected in a snapshot, the remainder of the snapshot optionally does not need to be scanned and the next snapshot in the snapshot chain can be searched for malware.

Further, the mounting 904 through the detecting 910 can be repeated for all objects of a system including virtual machines, file systems, databases, network attached storage, etc. When malware is detected, metadata in the snapshot can be adjusted to indicate quarantine status. For example, quarantine-related metadata for the infected snapshots (e.g., quarantine status metadata, among other quarantine-related metadata) may be updated for the infected snapshot. Additionally, or alternatively, a ledger (e.g., an append-only ledger, a snapshot-level ledger, etc.) used to keep track of snapshots that have been quarantined, released from quarantine, or both, may be maintained. In some examples, the quarantine related-metadata for a snapshot is updated by updating related fields in a snapshot-level quarantine metadata table. Quarantining a snapshot is described in more detail herein, including with reference to FIG. 11.

At operation 912, the GUI 524 generates an interface, such as the example interface 700. For each object, the example interface 700 can show snapshots over time with a status (e.g., color coded) of each snapshot (infected, not infected, encrypted, infected and partially encrypted). The example interface 700 may also illustrate a cut point above which snapshots are quarantined (e.g., not recoverable, recovery not permitted, recoverable with sufficient user privileges, partially recoverable, etc.). The GUI 524 may display object(s) scanned (VM, Host/Share/fileset, etc.); snapshot(s) scanned; Date/time of scan; scan filter criteria; Hash/Rules, etc. being scanned for (e.g., for YARA Rules: i. Rule name ii. Namespace iii. Tags iv. Hash of the rule, etc.); number of matched files; and number of matches; etc.

The GUI 524 may also provide detailed results after a search including: 1. The file name & path where a match was found 2. The rule (hash, or YARA name, namespace and tags) that were matched 3. The time when this file was

23

created 4. The time when this file was modified 5. The owner of this file a. Their fully qualified name (if retrievable) b. Any security identifier (e.g., SID for Windows, User ID for Linux/Unix) 6. The MD5, SHA1, and SHA256 hashes of this file, etc.

The GUI **524** then receives a command from a user to recover an infected snapshot at operation **914**. The privilege determinator **518** then determines at operation **916** if the user has sufficient privileges to recover an infected snapshot based on the user privileges **520**, which lists recovery privileges for users for infected snapshots and files. If the user has insufficient privileges, the recoverer **516** will not recover the infected snapshot and the method **900** ends. Otherwise, if the user has sufficient privileges at operation **916** then the recoverer **516** will recover (e.g., mount, restore, examine, read, etc.) the selected infected snapshot at operation **918**. Optionally, the selected snapshot can be recovered to a sandboxed virtual machine. The method **900** then ends.

In an embodiment the method **900** may further comprise propagating the quarantine status to other infected snapshots (e.g., replicas of infected snapshots). Quarantine status may be indicated in metadata of a snapshot (e.g., via setting a bit) and the related snapshots can be marked similarly by looking up a record of snapshots and replicas and then marking metadata of the related (e.g., replica) snapshot.

FIG. **10** depicts a flowchart illustrating a method **1000** of recovering non-infected content within an infected snapshot according to an example embodiment. In an example embodiment, the storage appliance **300** can execute the method **1000** using the malware engine **502**. Example methods described herein may also be implemented in the form of executable instructions stored on a machine-readable medium or in the form of electronic circuitry. For instance, the operations of the method **1000** may be represented by executable instructions that, when executed by a processor of a computing device, cause the computing device to perform the method **1000**. Depending on the embodiment, an operation of an example method described herein may be repeated in different ways or involve intervening operations not shown. Though the operations of example methods may be depicted and described in a certain order, the order in which the operations are performed may vary among embodiments, including performing certain operations in parallel.

At operation **1002**, the hypervisor **208** generates a virtual machine, e.g., virtual machine **220**, which can be sandboxed (e.g., no network access to prevent the spread of malware to other parts of a network). A snapshot is then mounted at operation **1004**, e.g., with Rubrik's LIVEMOUNT. The snapshot may be part of a snapshot chain and accordingly, a most recent snapshot may be mounted or a specific snapshot selected for mounting. Data within the snapshot is then hydrated with the hydrator **504** at operation **1006**, which can include deduplicating data. The detector **508** then searches the mounted hydrated snapshot for malware at operation **1008**. Searching can include applying YARA rules **510** and/or looking for matching hashes **512** (e.g., MD5, SHA1, and/or SHA256) and/or looking for specific file names (e.g., by name, such as ABAP/Rivpas.c.intd, and/or by path prefix) and/or path matches. The YARA rules **510** and/or hashes **512** may be imported (e.g., if new YARA rules **510** or hashes **512** are created) and the detector **508** can search for malware **1008** using all or a subset of the YARA rules **510** and/or hashes **512**.

For example, given a set of Indicators of Compromise encoded as YARA rules in a CISA alert (such as Alert (AA20-302A) Ransomware Activity Targeting the Health-

24

care and Public Health Sector), the detector **508** may identify any snapshots, file paths, or both, where those indicators were found. Or, given a set of suspicious file hashes or file paths (such as from Microsoft's Hafnium IOC list), the detector **508** may identify those snapshots, file paths, or both as well.

Further, the detector **508** may restrict a search to a specific snapshots or range of snapshots, a subset of files/directories and/or scan a replica of a snapshot instead an original snapshot. Restricting the search can include specific subsets of files or directories; file size limits, file creation/modification timestamps (e.g., before, after between timestamp (s)); file ownership; and files added/modified in a current snapshot (i.e., files that were not in prior snapshot or modified since that snapshot).

If, at operation **1010**, malware is detected, then a next most recent snapshot in the snapshot chain is mounted and the process is repeated until no malware is detected at operation **1010**. Further, in case of false negatives, additional snapshots in the chain can be scanned after no malware is detected in operation **1010**. Alternatively, the mounting **1004** through the detecting **1010** can start with an oldest or base snapshot and assuming no malware was detected in the base snapshot, repeating the process until malware is detected.

Further, the mounting **1004** through the detecting **1010** can be repeated for all objects of a system including virtual machines, file systems, databases, network attached storage, etc. When malware is detected, metadata in the snapshot can be adjusted to indicate quarantine status. For example, quarantine-related metadata for the infected snapshots (e.g., quarantine status metadata, among other quarantine-related metadata) may be updated for the infected snapshot. Additionally, or alternatively, a ledger (e.g., an append-only ledger, a snapshot-level ledger, etc.) used to keep track of snapshots that have been quarantined, released from quarantine, or both, may be maintained. In some examples, the quarantine related-metadata for a snapshot is updated by updating related fields in a snapshot-level quarantine metadata table. Quarantining a snapshot is described in more detail herein, including with reference to FIG. **11**.

At operation **1012**, the GUI **524** generates an interface, such as the example interface **700**. For each object, the example interface **700** can show snapshots over time with a status (e.g., color coded) of each snapshot (infected, not infected, encrypted, infected and partially encrypted). The example interface **700** may also illustrate a cut point above which snapshots are quarantined (e.g., not recoverable, recovery not permitted, recoverable with sufficient user privileges, partially recoverable, etc.). The GUI **524** may display object(s) scanned (VM, Host/Share/fileset, etc.); snapshot(s) scanned; Date/time of scan; scan filter criteria; Hash/Rules, etc. being scanned for (e.g., for YARA Rules: i. Rule name ii. Namespace iii. Tags iv. Hash of the rule, etc.); number of matched files; and number of matches; etc.

The GUI **524** may also provide detailed results after a search including: 1. The file name & path where a match was found 2. The rule (hash, or YARA name, namespace and tags) that were matched 3. The time when this file was created 4. The time when this file was modified 5. The owner of this file a. Their fully qualified name (if retrievable) b. Any security identifier (e.g., SID for Windows, User ID for Linux/Unix) 6. The MD5, SHA1, and SHA256 hashes of this file, etc.

At operation **1014**, the most recent non-infected snapshot is recovered (e.g., mounted, restored, etc.). Then, non-infected content related to the recovered snapshot is identi-

25

fied in more recent infected snapshots based on the prior searching **1008** at operation **1016**. Then content (e.g., files) from the non-infected snapshot selected by a user can be recovered in the more recent snapshots using forward incremental recovery until an infected or encrypted content is reached corresponding to the selected content. Alternatively, all content can be restored using forward incremental recovery for each content until infected or encrypted content is reached. In this way, the most recent non-infected content available is recovered even if some of the snapshots holding relevant content are infected.

For example, for a selected content including a file in a non-infected snapshot, that file can first be recovered in the most recent non-infected snapshot, then starting with the next (infected) snapshot in the snapshot chain, the next incremental file can be recovered, and repeated until infected or encrypted content is reached. Alternatively, as the snapshots have already been searched for malware at **1008**, the most recent infected snapshot with the selected file can be identified and the selected file recovered from that infected snapshot. In effect, while the example interface **700** shows a cut point at the snapshot level, there may be a cut point further up in the snapshot chain on a finer grained level (e.g., content level). The method **1000** then ends.

FIG. **11** depicts a flowchart illustrating a method **1100** of quarantining information in a snapshot according to an example embodiment. In an example embodiment, the storage appliance can execute the method **1100** using the flagger **514**. In some examples, the storage appliance is implemented at the data center (e.g., such as storage appliance **300**). In some examples, the storage appliance is implemented external to the data center (e.g., such as storage appliance **102**). When the storage appliance is implemented external to the data center, the functionality described herein may be provided using a software-as-a-service model. In such cases, the storage appliance may provide a quarantine functionality to end users. The storage applicant may provide access to a service that is used for quarantining information, identifying a range of information to be quarantined (e.g., based on identifying first information to be quarantined); determining a quarantined state of information (e.g., a snappable, snapshot, volume, section, folder, or file); and identifying whether information should be recover (e.g., based on whether the information or related information is quarantined), among other functions. The storage appliance may maintain a ledger that keeps track of the quarantined/released status of information (e.g., the user's information). Also, the storage appliance may provide a user interface for indicating quarantine statuses and releasing and recovering quarantined information.

Example methods described herein may also be implemented in the form of executable instructions stored on a machine-readable medium or in the form of electronic circuitry. For instance, the operations of the method **1100** may be represented by executable instructions that, when executed by a processor of a computing device, cause the computing device to perform the method **1100**. Depending on the embodiment, an operation of an example method described herein may be repeated in different ways or involve intervening operations not shown. Though the operations of example methods may be depicted and described in a certain order, the order in which the operations are performed may vary among embodiments, including performing certain operations in parallel.

At **1102**, a ledger may be initialized (e.g., by the flagger **514**) for recording quarantine-related operations—e.g., quarantine operations, release operations (which may also

26

be referred to as unquarantine operations). Initializing the ledger may include allocating memory to the ledger, clearing any data previously stored for the ledger, enabling the use of the ledger, and the like.

The ledger may be an append-only ledger. In such cases, after each quarantine-related operation performed, an indication of the latest performed operation may be added to an end of the ledger. The ledger may include a field for indicating an index of a recorded operation in the ledger, a field for indicating a time when a recorded operation was performed, a field for identifying one or more snapshots associated with a recorded operation, a field for identifying a file or folder that triggered the operation, a field for indicating a quarantine-related status of a snapshot, file, or folder associated with the recorded operation, or any combination thereof. In some examples, the ledger includes a field for identifying an object captured by a snapshot (e.g., a snappable) and indicating a time a snapshot was taken. In some examples, the ledger may include a field identifying a user (e.g., an internal system user or external user) that triggered an operation, a field for user notes (e.g., a description as to why an operation was performed), and a field for operation notes (e.g., a description of a performed operation or result of the performed operation), or any combination thereof. A ledger for recording quarantine-related operations is described in more detail herein, including with reference to FIG. **12**.

In some examples, the ledger may be an operation-level ledger. In such cases, a record of each quarantine-related operation (e.g., a quarantine operation, release operation, etc.) performed by the storage appliance for a target object may be maintained. Thus, an operation-level ledger may be used to determine a quarantine state of information in one or more snapshots of the target object as well as to record a history of quarantine operations performed for the target object.

In some examples, the ledger may be a snapshot-level ledger. In such cases, a record of whether a snapshot is in an accessible state, a quarantined state, or a released state may be maintained. In some examples, a record for a snapshot may not be added to the snapshot-level ledger until the snapshot (or information within the snapshot) has been quarantined—e.g., the snapshot-level ledger may have a sparse database structure. In some examples, a record for a snapshot that has been quarantined may be removed from the snapshot-level ledger when the snapshot (or the information within the snapshot) is released from quarantine. In such cases, if a snapshot is not included in the snapshot-level ledger, an absence of the snapshot in the ledger may indicate that the snapshot is in an accessible state. In other examples, each recorded snapshot may be added to the snapshot-level ledger along with a current status (e.g., accessible, quarantined, released) of the snapshot. In some examples, when a snapshot-level ledger is used, the storage appliance may generate a separate audit log for each quarantine-related operation performed to record a history of quarantine operations for a target object. A snapshot-level ledger for recording quarantine-related operations is described in more detail herein, including with reference to FIG. **15**.

At **1104**, snapshots of one or more target objects may be stored (e.g., in storage device **104** by the storage appliance **300**). A target object may be a virtual machine, fileserver, workstation, network attached storage, database, etc. In some examples, a snapshot includes a complete image of a target object (such a snapshot may, for example, be referred to as a base snapshot for the target object). A snapshot may also include an incremental or differential image of a target

object—e.g., relative to a complete image of the target object (such a snapshot may, for example, be referred to as an incremental snapshot for the target object). In some examples, snapshots of the one or more target objects may be taken in accordance with a schedule, whenever a change is detected at the target object, or both. For example, a snapshot may be taken of a virtual machine every hour and a snapshot of a file system may be taken each time information in the file system changes.

At 1106, an indication of information flagged for quarantine may be obtained (e.g., at the storage appliance 300 or the storage appliance 102). In some examples, an indication of one or more files flagged for quarantine may be received from the malware engine 502—e.g., based on the malware engine 502 detecting malware within one or more snapshots of a target object. In some examples, the malware engine 502 detects malware within a target object based on detecting an infected or suspicious file or based on detecting maliciously encrypted files in a snapshot. Additionally, or alternatively, an indication of one or more files flagged for quarantine may be received from a user associated with (e.g., that manages) a target object (e.g., a system administrator). In some examples, the user detects and identifies infected, suspicious, or maliciously encrypted files during operation of the target object and sends an indication of the affected files to the storage appliance as being flagged for quarantine. Additionally, or alternatively, the user may detect files that cause or are suspected as causing system crashes, are security vulnerabilities, are not ready for release, and the like, and may send an indication of the identified files to the storage application as being flagged for quarantine. In some examples, the indication sent by the malware engine 502 may include an indication of a snapshot that includes the flagged files.

At 1108, the information flagged for quarantine may be identified (e.g., by the storage appliance 300 or the storage appliance 102). In some examples, the storage appliance compares indicated files with files included in one or more snapshots stored for a target object. For example, the storage appliance may compare the indicated files with each file in each snapshot stored for a target object or for a subset of snapshots stored for the target object (e.g., snapshots taken within a time range). In some examples, the storage appliance uses the hydrator 504 and mounter 506 to materialize and mount the snapshots of the target object (e.g., on a snapshot-by-snapshot basis). The storage appliance may further use the flagger 514 to identify and flag snapshots including the indicated files as well as a file path of the indicated files.

In some examples, the storage appliance may further use the flagger 514 and propagator 522 to identify additional snapshots associated with a snapshot for a target object that includes flagged information. In some examples, the storage appliance identifies earlier versions of a snapshot (e.g., prior snapshots used for restoring the snapshot) that include information flagged for quarantine based on determining that a snapshot includes the information flagged for quarantine. Flagger 514 may flag the identified snapshots as quarantined.

In some examples, the propagator 522 may be used to identify snapshots for the target object occurring after the snapshot, that are based on the snapshot, or both, and flagger 514 may flag the identified snapshots as quarantined. In some examples, before flagging the identified snapshots, the storage appliance may determine whether the identified snapshots include the flagged information. In some cases, after identifying quarantined information in a snapshot, the

propagator 522 identifies additional snapshots (e.g., earlier or later snapshots) in one or more chains of snapshots that includes the snapshot. In some examples, an object may be recovered using the snapshots in a chain of snapshots, where different snapshots in the chain of snapshots include partial information for recovering a full set of information.

In some cases, after identifying quarantined information in a snapshot for an object, the flagger 514 may flag information in subsequent snapshots of the object—that occur within a time range of the snapshot. The subsequent snapshots may include snapshots for the object that are outside of a snapshot chain including the initial snapshot. In some examples, the time range is indicated by a user. For example, the flagger 514 may flag subsequent snapshots taken of the object until the user instructs the flagger 514 to stop. In other examples, the user may indicate to the flagger 514 to flag subsequent snapshots that occur up to a designated time period (e.g., a month) after the initial snapshot was taken.

In some examples, the storage appliance identifies versions of the snapshot stored at other locations (e.g., at edge servers, cloud servers, or on-site customer servers). Flagger 514 may flag the versions of the snapshot stored at the other locations as quarantined.

At 1110, the ledger for keeping track of quarantine-related operations (which may also be referred to as a “quarantine metadata table”) may be updated based on the identified snapshots, files, or both (e.g., using flagger 514). For example, for each snapshot identified as including flagged information, an entry may be added to the ledger that identifies the snapshot and indicates that the snapshot is in a quarantined state. Additionally, or alternatively, for each file or folder identified as including flagged information, an entry may be added to the ledger that identifies the file or folder and indicates that the file or folder is in a quarantined state. Each entry added to the ledger may also include an indication of a position of the entry within the ledger (e.g., using an index), a time at which the associated information was created (e.g., a time when a snapshot was taken), a time at which the associated information flagged for quarantine was detected, a time at which the entry was added to the ledger, or any combination thereof. In some examples, entries in the ledger that exceed a threshold age (e.g., two years) may be removed from the ledger.

In some examples, if a snapshot-level ledger is used, updating the ledger may include adding a snapshot to the ledger that has been flagged for quarantine or that contains information that has been flagged for quarantine. In other examples, a set of snapshots may already be indicated by the snapshot-level ledger, and updating the snapshot-level ledger may include changing a status of the snapshot (e.g., from an accessible or released state to a quarantined state) and updating the files of the snapshot indicated as quarantined (e.g., if other files of the snapshot are already quarantined).

At 1112, an indication of information flagged for release from quarantine may be received. In some examples, the indication may indicate one or more files, folders, or snapshots to release from quarantine. In some cases, the indication indicates one or more files or folders of a target object to release from quarantine, and the storage appliance identifies a set of snapshots for the target object that includes the files or folders. The indication may be received from a user having certain privileges.

At 1114, an operation for identifying information to release from quarantine may be performed (e.g., the storage appliance 300 or the storage appliance 102). In some examples, the operation is performed after determining that

the user associated with the indication has a specific or a threshold level of privileges. For example, the storage appliance may perform the operation if the user is a specific user, has administrative privileges, has super-administrative privileges, or a combination thereof. In some examples, before performing the operation, the storage appliance may authenticate the user by requesting a password or key from the user. In some examples, the password or key is different than a login password for the user—e.g., the password may be specific to enabling release operations.

In some examples, as part of the operation for identifying information to release, the storage appliance may identify one or more snapshots including the identified information, as similarly described with reference to identifying information flagged for quarantine.

At **1116**, the ledger may be updated (e.g., by flagger **514**) based on identifying information to release from the quarantine. Updating the ledger may include adding, for each file path, folder path, or snapshot identified as including the indicated information, an entry identifying the file path, folder path, or snapshot and indicating a released state of the file path, folder path, or snapshot.

In some examples, the indication of information flagged for release identifies a file path, folder path, snapshot, or any combination thereof associated with the information. In such cases, instead of identifying information to release in stored snapshots, the flagger **514** may directly add an entry to the ledger including the indicated file path, folder path, or snapshot identifier and a corresponding released state.

In some examples, when a snapshot-level ledger is used, updating the ledger may include removing or updating a record for a snapshot including the information released from quarantine. Updating the record may include updating a field in the record that indicates information in the snapshot that is quarantined so that the field indicates information that is currently quarantined. In some examples, the record for the snapshot may be removed if all of the information in the snapshot that was quarantined is identified as being released from quarantine.

In some examples, the ledger includes each snapshot regardless of quarantine status. In such cases, the record for the snapshot may not be removed; instead, a status indicated in a field of the record may be updated. For example, if all of the information in the snapshot that was quarantined is identified as being released from quarantine, a status field of the record for the snapshot may be updated to indicate that the record is accessible or has been released from quarantine. Alternatively, if a portion of the information in the snapshot that was quarantined is identified as being released from quarantine, the status field of the record may be maintained and a quarantined information identifier field may be updated to remove the released information.

FIG. 12 illustrates an example ledger that supports recording quarantining and release operations in accordance with examples as disclosed herein.

The ledger **1200** is configured to keep track of snapshots, files, folders, or any combination thereof, that are in a quarantine state as well as snapshots, files, folders, or any combination thereof, that have been released from a quarantine state. In some examples, a ledger (such as the ledger **1200**) is maintained for each object backed up by a storage service. The ledger **1200** may include an index field **1205**, a time field **1210**, a snapshot identifier field **1215**, a path identifier field **1220**, and an operation identifier field **1225**.

The index field **1205** may be used to indicate a position of an entry within the ledger **1200**. In some examples, the

indices may be used to determine a relative timing of the quarantining or release of information indicated by different entries.

The time field **1210** may be used to indicate a time when the associated information was quarantined or released, when the entry was added to the ledger, or both. Similar to the indices, the timing may similarly be used to determine a relative timing of the quarantining or release of information indicated by different entries.

The snapshot identifier field **1215** may be used to indicate a snapshot that was identified as including information flagged for quarantining or release from quarantine. In some examples, the snapshot identifier field **1215** may indicate multiple snapshots flagged for quarantining or release from quarantine.

The path identifier field **1220** may be used to indicate a path to a file, folder, section, or volume that was identified as including information flagged for quarantining or release from quarantine. The indicated path may be within the snapshot or snapshot(s) identified by snapshot identifier field **1215**.

By including the path identifier field **1220** in the ledger **1200**, the ledger **1200** may be used to support file-level quarantining. Also, the path identifier field **1220** may be used to indicate to a user which information is preventing restoration. In some examples, the ledger **1200** omits the path identifier field **1220**—e.g., if quarantining is performed solely on a snapshot level.

The operation identifier field **1225** may be used to indicate a type of operation performed on the information associated with an entry—e.g., whether the operation was a quarantine operation or a release operation. In some examples, the operation identifier field **1225** indicates that information associated with an entry is quarantined. In other examples, the operation identifier field **1225** indicates that information associated with an entry has been released from quarantine.

In some examples, the ledger **1200** may also include a snapshot time field that indicates a time when a snapshot was taken. The snapshot time field may also include an identifier of an object captured by a snapshot (e.g., a snappable). In some examples, the snapshot time field may be included in the snapshot identifier field **1215**. In some examples, the snapshot identifier and snapshot time may be used to identify all quarantine operations performed for an object in a given time range. In some examples, whether a point-in-time copy of an object is quarantined may be based on identifying the quarantine operations performed on an object during a time range—e.g., if the point-in-time copy was obtained within the time range. In some examples, continuous data protection software may use the snapshot identifier and snapshot time information to determine whether a point-in-time copy is quarantined.

In some examples, the snapshot time and the identifier of an object captured by the snapshot may be included in separate fields (e.g., the snapshot time field and the snapshot identifier field **1215**). In either case, the identifier of the object may include an index, which may be used to improve a performance of queries by facilitating the identification of all ledger entries for a given object.

The ledger **1200** may also include a hash field. The hash field may include a hash of information (e.g., a section, folder, or a file) flagged for quarantined. The hash may be used to determine whether the information flagged for quarantine has changed since being quarantined. In some examples, a decision of whether to release information from quarantine may be made based on whether the information has changed. For example, if quarantined information has

changed since being quarantined, the information may not be recovered. Additionally, or alternatively, the hash may be used to ensure that the file remains quarantined in subsequent snapshots even if the path of the information changes.

In some examples, the ledger **1200** may include a user account field, a user notes field, an operation notes field, or any combination thereof. A user account field may indicate a user associated with a performed operation (e.g., a quarantine or release operation). For example, the user that triggered the operation may be indicated in the user account field. In some examples, the user is an internal system user—e.g., a system user for a threat analysis application. In some examples, the user is an external user (e.g., an IT or system administrator).

A user notes field may be used to indicate notes associated with a performed operation. In some example, the user notes field may indicate why the operation was performed. For example, the user notes field may indicate that a quarantine operation was performed because a suspicious file was identified or unexpectedly encrypted data was identified. In some examples, the user notes field may indicate that a quarantine operation was performed because a system administrator had concerns that particular information was creating system failures. In yet other examples, the user notes field may indicate that a release operation was performed because the information identified by the threat analysis application was a false positive or based on identifying that the quarantined information was not causing the system failures.

An operation notes field may provide information related to the operation. For examples, the operation notes field may indicate that the operation succeeded, a time when the operation was performed, or the like.

In some examples, when information flagged for quarantine is identified, one or more entries associated with the information may be added to the ledger **1200**. And if quarantined information is flagged for release, the entry may be removed from the ledger **1200**.

In other examples, the ledger **1200** may be an append-only ledger. For example, instead of removing an entry associated with quarantined data if the information associated with the entry is released, an additional entry associated with the information may be added at an end of the ledger **1200** to indicate that the information has been released from quarantine. That is, the ledger **1200** may include a field indicating a quarantine action associated with an entry—e.g., operation identifier field **1225**. For example, if information is placed into quarantine, a corresponding entry of the ledger **1200** may indicate a quarantine action for the information. And if the information is released from quarantine, another corresponding entry of the ledger **1200** may indicate a quarantine release action for the information. The ledger **1200** may at any time be consulted and analyzed (e.g., by a storage appliance **300** or a storage appliance **102**) to determine the quarantine status of a piece of information (e.g., snapshot) based on the net effect of all entries associated with the information, considering the chronological order of all such entries. For example, if an entry indicating a quarantine action for the information is identified within the ledger **1200**, and no subsequent entry indicating a quarantine release action for the information is identified, then it may be determined that the information is quarantined. Alternatively, if an entry indicating a quarantine action for the information is identified within the ledger **1200**, a subsequent entry indicating a quarantine release action for the information is identified, and no further subsequent entry

indicating another quarantine action for the information is identified, then it may be determined that the information is not quarantined.

In some examples, an indication of information flagged for quarantine is received (e.g., at a storage appliance **300** or a storage appliance **102**). The indication may also include an indication of the target object including the information. The storage appliance may analyze a set of snapshots of the target object to determine whether any of the snapshots include the information flagged for quarantine. Based on the analysis, the storage appliance may identify one or more snapshots that include the information—e.g., based on matching one or more indicated files or folders with corresponding files or folders in the one or more snapshots.

In some examples, the storage appliance may add a first entry to the ledger **1200** based on identifying a snapshot that includes the flagged information. The entry may include an index (e.g., index 1), a time when the entry was added to the ledger **1200** and/or when the flagged information was detected, an identifier of the snapshot (e.g., snapshot A), an identifier of the path (e.g., file path, folder path, section ID) to the flagged information, and an indication that the snapshot, file, folder, or any combination thereof, is quarantined. Based on identifying a second snapshot that includes the flagged information, the storage appliance may similarly add a second entry to the ledger **1200**—e.g., with an index 2 and an identifier of a snapshot C. In some examples, the storage appliance may identify additional snapshots (e.g., snapshot D) that depend from the identified snapshots (e.g., snapshot A or B) and similarly add additional entries to the ledger **1200**. In some examples, instead of adding additional entries to the ledger, the storage appliance may include an indicator of the snapshots depending from an identified snapshot in the snapshot identifier field **1215** of the entry.

In some examples, the storage appliance may add additional entries to the ledger **1200** based on receiving an indication that additional information within a quarantined snapshot is flagged for quarantine. For example, the storage appliance may receive an indication that additional files in snapshot C have been flagged for quarantine. The storage appliance may proceed to add the fifth entry to the ledger **1200** indicating that the additional files in snapshot C are quarantined.

Later, a subsequent indication that information has been flagged for release from quarantine may be received. The storage appliance may analyze a set of snapshots of the target object to determine whether any of the snapshots include the information flagged for release. Based on the analysis, the storage appliance may identify one or more snapshots that include the information—e.g., based on matching one or more indicated files or folders with corresponding files or folders in the one or more snapshots. Alternatively, the storage appliance may analyze the ledger **1200** to determine if any of the quarantined snapshots include the information (e.g., the file path, folder path, or section ID) flagged for release.

In some examples, the indication indicates that a snapshot has been released from quarantined. In some examples, the indication indicates that a section, folder, or file has been released from quarantine.

The storage appliance may add one or more entries to the ledger **1200** indicating the release of quarantined information. For example, an entry may be added that includes an index (e.g., index 10), a time when the entry was added to the ledger **1200**, an identifier of the snapshot (e.g., snapshot A), an identifier of the path to the released information, and an indication that the snapshot, file, folder, or any combi-

nation thereof, has been released from quarantine (e.g., an indication of an updated status of the information, an indication of the quarantine release action for the information, or both).

In some examples, an entry may indicate that a section, folder, or file in a snapshot has been released from quarantine. For instance, the entry at index 12 may indicate that a folder in snapshot C has been released from quarantine. In some examples, an entry may indicate that an entire snapshot has been released. For instance, the entry at index 13 may indicate that snapshot E has been released from quarantine.

In some examples, the ledger **1200** includes empty entries that have not yet been filled—e.g., if the ledger **1200** is initialized with a certain size.

In some examples, the storage appliance may perform an operation for listing information that is quarantined for an object (which may be referred to as a list quarantined files operation). The list quarantined files operation may add an entry to the ledger **1200** that indicates information (e.g., a list of files or folders) that are quarantined in a snapshot. In some cases, a list of quarantined files may be based on previous quarantine and release operations performed on a snapshot (or earlier snapshots in a chain of snapshots). In some examples, a list of quarantined information for a snapshot may omit information that were listed as quarantined for an earlier snapshot—e.g., if some of the earlier information was released from quarantine or deleted between snapshots. In some examples, a list quarantined files operation may be performed for one or more snapshots, and the resulting list of quarantined files may be input as YARA rules while searching for malware.

FIG. 13 illustrates an example ledger that supports recording quarantining and release operations in accordance with examples as disclosed herein.

The ledger **1300** is configured to keep track of information (e.g., snapshots, files, folders) that are in a quarantine state. In some examples, the ledger **1300** may also be configured to keep track of information that is in an accessible state (e.g., implicitly or explicitly). The ledger may also be configured to keep track of information that is in a released state.

The ledger **1300** may include a snapshot identifier field **1315**, a time field **1310**, a path identifier field **1320**, which may be respective examples of the snapshot identifier field **1215**, the time field **1210**, the path identifier field **1220** of FIG. 12. The ledger **1300** may also include a quarantine status field **1325**. In some examples, the ledger **1300** may include additional fields—e.g., any one or more of the fields supported by the ledger **1200** of FIG. 12 as described herein including with reference to FIG. 12. For example, among others, the ledger **1300** may include an index field, snapshot time field, a hash field, a user account field, a user notes field, an operation notes field, or any combination thereof.

Quarantine status field **1325** may indicate that the information in a corresponding snapshot identified in the snapshot identifier field **1315** is quarantined. In some examples, quarantine status field **1325** is omitted from ledger **1300**—e.g., if it is implied that any information included in the ledger **1300** is quarantined.

In some examples, an indication of information flagged for quarantine is received (e.g., at a storage appliance **300** or a storage appliance **102**). The indication may also include an indication of the target object including the information. The storage appliance may analyze a set of snapshots of the target object to determine whether any of the snapshots include the information flagged for quarantine. Based on the

analysis, the storage appliance may identify one or more snapshots that include the information—e.g., based on matching one or more indicated files or folders with corresponding files or folders in the one or more snapshots.

In some examples, the storage appliance may add a first entry to the ledger **1300** based on identifying a snapshot that includes the flagged information. The entry may include an identifier of the snapshot (e.g., snapshot A), a time when the entry was added to the ledger **1300** and/or when the flagged information was detected, an identifier of the path to the flagged information, and an indication the identified information is quarantined. The storage appliance may similarly add a second entry to the ledger **1300** based on identifying that a second snapshot includes flagged information.

In some examples, the storage appliance adds entries for snapshots that are not identified as including quarantined information. In such cases, the quarantine status field **1325** may indicate that the snapshots are accessible. In other examples, entries for unquarantined snapshots may be omitted from the ledger **1300**. In such cases, adding an entry in the ledger may imply that the snapshot indicated by the snapshot identifier field **1315** and/or the information indicated by the path identifier field **1320** is quarantined.

Later, a subsequent indication that additional information has been flagged for quarantine may be received. In some examples, the storage appliance may determine that the additional information is included in a snapshot that already includes quarantined information (e.g., snapshot A). In such cases, the storage appliance may update the entry for snapshot A—e.g., by updating the path identifier field **1320** to indicate the additional information as well as the previously quarantined information.

Later, a subsequent indication that information has been flagged for release from quarantine may be received. In some examples, the information released from quarantine is included in snapshot A, where the released information may be a subset of the quarantined information in snapshot A. In such cases, the path identifier field **1320** of the first entry may be updated to no longer indicate the released information. In some examples, the information released from quarantine is included in snapshot B, where the released information may encompass the quarantined information in snapshot B. In such cases, the second entry may be removed from the ledger **1300**. In other examples, instead of removing the second entry, the quarantine status field **1325** of the second entry may be updated to indicate that snapshot B has been released from quarantine.

FIG. 14 depicts a flowchart illustrating a method **1400** of restoring a requested snapshot in accordance with quarantine information according to an example embodiment. In an example embodiment, the storage appliance **300** can execute the method **1400** using the recoverer **516**. In some examples, the storage appliance is implemented at the data center (e.g., such as storage appliance **300**). In some examples, the storage appliance is implemented external to the data center (e.g., such as storage appliance **102**).

Example methods described herein may also be implemented in the form of executable instructions stored on a machine-readable medium or in the form of electronic circuitry. For instance, the operations of the method **1400** may be represented by executable instructions that, when executed by a processor of a computing device, cause the computing device to perform the method **1400**. Depending on the embodiment, an operation of an example method described herein may be repeated in different ways or involve intervening operations not shown. Though the operations of example methods may be depicted and

described in a certain order, the order in which the operations are performed may vary among embodiments, including performing certain operations in parallel.

At **1402**, a request to restore information may be received (e.g., at the storage appliance **300** or a storage appliance **102**). The request to restore information may include an indication of a file, folder, section, volume, or snapshot associated with a target object.

In some examples, the request may be intercepted by a software layer (which may be referred to as a QuarantineEnforcer) that is used to enforce quarantine conditions for a restoration procedure of a requested snapshot. In some examples, the software layer may be implemented (e.g., as middle-ware) between a front-end of an API that provides access to an application for restoring a requested snapshot and the application for restoring a requested snapshot. The software layer may intercept each call to the API, including indirect calls to the API received from a user via an interface of a system that contains the application and direct calls to the application via a user-accessible interface. The QuarantineEnforcer may also determine a privilege of a user based on the call to the API.

At **1404**, a chain of backup information for restoring the requested information may be identified (e.g., by the storage appliance **300** or a storage appliance **102**). For example, if the request is for a snapshot of a target object, the storage appliance may identify one or more additional snapshots for rebuilding the requested snapshot—e.g., if the requested snapshot is an incremental or differential snapshot. Similarly, if the request is for a file in a target object, the storage appliance may identify one or more additional files for rebuilding the file. For example, if a request to restore a snapshot **C** is received, the storage appliance may determine that snapshots **A** and **B** are needed to reconstruct snapshot **C**. In another example, if a request to restore a snapshot **G** is received, the storage appliance may determine that snapshots **C** is also are needed to reconstruct snapshot **G**.

In some examples, a sequence of backup information for an object may be identified. For example, if the request is for a snapshot of a target object, the storage appliance may identify one or more additional snapshots taken for the target object—e.g., the storage appliance may identify one or more snapshots taken within a time range of the requested snapshot. For instance, the storage appliance may identify snapshots taken up to a month before the snapshot, snapshots taken up to a month after the snapshot, or snapshots taken a half-month before and a half-month after the snapshot. Similarly, if the request is for a file in a target object at a time point, the storage appliance may identify one or more additional versions of the file that occurred within a time range of the time point.

At **1406**, a ledger (e.g., the ledger **1200** or the ledger **1300**) may be consulted (e.g., by the storage appliance **300** or a storage appliance **102**) to determine whether the request to restore the information is to be granted. When the ledger **1200** is used, consulting the ledger may include analyzing each entry in the ledger to determine whether any of the information used to reconstruct the requested information has been quarantined. Consulting the ledger may also include analyzing each entry in the ledger to determine whether any of the information has been released from quarantine. When the ledger **1300** is used, consulting the ledger may include identifying and analyzing an entry for one or more snapshots including the information requested for restoration.

At **1408**, whether any information used for reconstructing the requested information is in a quarantined state may be

determined (e.g., by the storage appliance **300** or a storage appliance **102**) based on consulting the ledger.

In some examples and with reference to the ledger **1200**, if a request to restore snapshot **A** is received, the storage appliance may determine that snapshot **A**, a file in snapshot **A**, or both was quarantined based on the first entry of the ledger **1200**. The storage appliance may further determine that snapshot **A**, the file in snapshot **A**, or both has been released from quarantine based on the tenth entry of the ledger **1200**—e.g., based on the tenth entry occurring after the first entry, based on the time entries, etc.

In some examples and with reference to the ledger **1200**, if a request to restore snapshot **C** is received, the storage appliance may determine that snapshot **C**, a folder in snapshot **C**, a file in snapshot **C**, or any combination thereof, was quarantined based on the second entry and fifth entry of the ledger **1200**. The storage appliance may further determine that the folder in snapshot **C** was released from quarantine based on the twelfth entry of the ledger **1200**. However, the storage appliance may determine that the file in snapshot **C** remains quarantined based on the information identified by the twelfth entry not including the information identified by the fifth entry.

In some examples and with reference to the ledger **1200**, if a request to restore snapshot **D** is received, the storage appliance may determine that snapshot **D**, a file within snapshot **D**, or both, was quarantined based on the fourth entry of the ledger **1200**. The storage appliance may further determine that snapshot **D**, the file within snapshot **D**, or both, has been released from quarantine based on the eleventh entry of the ledger **1200**.

In some examples and with reference to the ledger **1200**, if a request to restore snapshot **E** is received, the storage appliance may determine that snapshot **E**, a first file in snapshot **E**, a second file in snapshot **E**, or both was quarantined based on the third entry of the ledger **1200**, the sixth entry of the ledger, or both. The storage appliance may further determine that all of the information in snapshot **E** has been released from quarantine based on the thirteenth entry of the ledger **1200**. In some examples, the storage appliance determines that all of the information in snapshot **E** has been released from quarantine based on determining that the latest entry for snapshot **E** (e.g., the thirteenth entry) includes an indication that all of the information in snapshot **E** has been released from quarantine.

In some examples and with reference to the ledger **1200**, if a request to restore snapshot **F** is received, the storage appliance may analyze the ledger **1200** to determine whether snapshot **F** is quarantined. Based on the analysis of the ledger **1200**, the storage appliance may determine that snapshot **F** is quarantined based on the Nth entry of the ledger **1200**.

In some examples and with reference to the ledger **1200**, if a request to restore a snapshot **G** is received, the storage appliance may determine that snapshot **G** is not quarantined based on identifying no entries identifying snapshot **G**. In some examples, before determining that snapshot **G** is not quarantined, the storage appliance may identify a set of snapshots (including snapshot **G**) for restoring a computing system to a time **G** (which may be referred to as the snapshot chain for snapshot **G**). In some cases, the storage appliance may determine that a snapshot included in the snapshot chain includes quarantined information (e.g., if the snapshot chain includes snapshot **C**).

In some examples and with reference to the ledger **1300**, if a request to restore snapshot **A** is received, the storage

appliance may determine that snapshot A, a file in snapshot A, or both was quarantined based on the first entry of the ledger 1300.

In some examples and with reference to the ledger 1300, if a request to restore snapshot B is received, the storage appliance may determine that the information in snapshot B is accessible for a restoration operation—e.g., based on failing to identify snapshot B in the ledger 1300 in cases where snapshot B was removed from the ledger or determining from the ledger 1300 that snapshot B is released from quarantine in cases where snapshot B is retained in the ledger.

In some examples and with reference to the ledger 1300, if a request to restore snapshot C is received, the storage appliance may determine that the information in snapshot C is accessible for a restoration operation—e.g., based on failing to identify snapshot C in the ledger 1300 or determining from the ledger 1300 that snapshot C is accessible in cases where the quarantine of all snapshots is reflected in the ledger 1300. In some examples, a procedure for consulting a snapshot-level ledger to determine whether requested information is in a quarantined state may be completed in less time than a procedure for consulting an operation-level ledger. Because, for example, each entry in an operation-level ledger may be consulted while a single entry in a snapshot-level ledger may be consulted.

In some examples, the QuarantineEnforcer determines whether a call to the API is attempting to recover quarantined information—e.g., by consulting a ledger, such as ledger 1200 of FIG. 12 or the ledger 1500 of FIG. 15. The Quarantine Enforcer may also determine a privilege of a user that sent the call to the API. Based on the determination of whether the call is attempting to recover quarantined information, the privilege of the user, or both, the QuarantineEnforcer may determine whether to block the call from reaching the application or allowing the call to proceed to the application. In some examples, the QuarantineEnforcer allows the call to reach the application based on determining that the call is not attempting to recover quarantined information, and the storage appliance may proceed to perform the operations described with reference to 1410. In some examples, the QuarantineEnforcer allows the call to reach the application after determining that the call is attempting to recover quarantined information based on determining that the user associated with the call has sufficient privileges to bypass quarantine flags, and the storage appliance may proceed to perform the operations described with reference to 1420. In cases where QuarantineEnforcer allows the call to reach the application after determining that the call is attempting to recover quarantined information based on determining that the user associated with the call has sufficient privileges, QuarantineEnforcer may perform aspects of the operations described with reference to 1412 through 1418.

In some examples, the QuarantineEnforcer blocks the call from reaching the application based on determining that the call is attempting to recover quarantined information. In such cases, the storage appliance may proceed to perform the operations described with reference to 1412.

At 1410, a restoration of the requested information may be initiated (e.g., by the storage appliance 300 or a storage appliance 102). In some examples, the restoration is initiated if the storage appliance determines that the requested information is not quarantined. For example, for a request to restore a snapshot, the storage appliance may initiate the restoration based on determining that the snapshot does not include quarantined information. In another example, for a

request to restore a section, folder, or file in a snapshot, the storage appliance may initiate the restoration based on determining that the section, folder, or file does not include quarantined information.

By initiating restoration of requested information based on determining that the requested information is not quarantined, the storage appliance may restore the requested information with less latency than if the storage appliance were to further determine whether earlier information used to rebuild the requested information is quarantined.

In some examples, the restoration is initiated if the storage appliance determines that none of the information used to restore the requested information is identified as being in a quarantined state. For example, for a request to restore a snapshot, the storage appliance may initiate the restoration based on determining that the snapshot as well as any other previously saved snapshots used for the restoration of the snapshot do not include quarantined information. In another example, for a request to restore a section, folder, or file in a snapshot, the storage appliance may initiate the restoration based on determining that any previously saved sections, folders, or files used for the restoration of the section, folder, or file does not include quarantined information.

By determining whether any preceding information for restoring requested information is quarantined before initiating restoration for the requested information, the storage appliance may reduce the risk of restoring unwanted (e.g., infected) information.

At 1412, a restoration of the requested information may be prohibited (e.g., by the storage appliance 300 or a storage appliance 102). In some examples, the restoration is prohibited if the storage appliance determines that the requested information is quarantined. For example, for a request to restore a snapshot, the storage appliance may prohibit the restoration based on determining that the snapshot includes quarantined information. In another example, for a request to restore a section, folder, or file in a snapshot, the storage appliance may prohibit the restoration based on determining that the section, folder, or file includes quarantined information.

In some examples, the restoration is prohibited if the storage appliance determines that any of the information used to restore the requested information is quarantined. For example, for a request to restore a snapshot, the storage appliance may prohibit the restoration based on determining that the snapshot as well as any other previously saved snapshots used for the restoration of the snapshot include quarantined information. In another example, for a request to restore a section, folder, or file in a snapshot, the storage appliance may prohibit the restoration based on determining that any previously saved sections, folders, or files used for the restoration of the section, folder, or file include quarantined information.

In some examples, image level recovery using the snapshot is prohibited if the storage appliance determines that a snapshot used for restoration of the requested information includes any information that is quarantined. In some examples, image level recovery and/or all file recovery using the snapshot are prohibited if the storage appliance determines that a snapshot used for restoration of the requested information includes any information that is quarantined. In some examples, image level recovery and/or recovery of quarantined information (e.g., sections, folder, or files) using the snapshot are prohibited if the storage appliance determines that a snapshot used for restoration of the requested information includes any information that is quarantined.

At **1414**, a message indicating that the restoration of the requested information is prohibited may be displayed (e.g., by the storage appliance **300** or a storage appliance **102**). The message may include an indication of a snapshot or information within a snapshot that has been quarantined. In some examples, the message includes a command for overriding the prohibition on the restoration of the requested information.

At **1416**, a command to override the prohibition may be received (e.g., at the storage appliance **300** or a storage appliance **102**). The command may include an indication of a user that sent the command. In some examples, the command includes a password for overriding the prohibition.

At **1418**, privileges of the user that sent the command may be determined (e.g., by the storage appliance **300** or a storage appliance **102**). In some examples, the storage appliance determines that the user is on a list of one or more users that are permitted to override the prohibition. In some examples, the storage appliance determines that the user is of a type (e.g., an administrator or super-administrator) that is permitted to override the prohibition. In some examples, after identifying that the user is permitted to override the prohibition, the storage appliance further sends a request for a password for overriding the prohibition.

At **1420**, restoration of the requested information may be initiated (e.g., by the storage appliance). In some examples, the requested information is restored to a sandboxed environment (e.g., a sandboxed virtual machine) despite the requested information including quarantined information. In some examples, the requested information is restored to a user's environment. In some examples, the restoration of the requested information is performed after determining that the user has sufficient privileges, receiving an override password, or both.

At **1422**, a command to perform a partial restoration may be received (e.g., at the storage appliance **300** or a storage appliance **102**) in response to the message indicating that the restoration of the requested information is prohibited.

At **1424**, the partial restoration may be initiated (e.g., by the storage appliance **300** or a storage appliance **102**). In some examples, the partial restoration restores the requested information except that any portion of the requested information that is quarantined is not restored. For example, if a request to partially restore a snapshot is received, the snapshot may be restored except that a set of files, folders, or sections used to reconstruct the snapshot that are quarantined may not be restored with the snapshot.

FIG. **15** illustrates an example ledger that supports recording quarantining and release operations in accordance with examples as disclosed herein.

The ledger **1500** is configured to keep track of snapshots, files, folders, or any combination thereof, that are in a quarantine state as well as snapshots, files, folders, or any combination thereof, that have been released from a quarantine state. In some examples, a ledger (such as the ledger **1500**) is maintained for each object (e.g., target) backed up by a storage service. The ledger **1500** may be referred to as a snapshot-level ledger. The ledger **1500** may include a snapshot identifier field **1515**, a path identifier field **1520**, and an operation identifier field **1525**, which may be examples of a snapshot identifier field **1215**, a path identifier field **1220**, and an operation identifier field **1225** described with reference to FIG. **12**. In some examples, the ledger **1500** may include additional or alternative fields, such as the additional field described with reference to the ledger **1200** of FIG. **12**.

In some examples, the ledger **1500** may include an entry for each snapshot captured by the data management system. Alternatively, the ledger **1500** may include an entry only for snapshots that have been subjected to a quarantine operation. In some examples, if the ledger **1500** includes an entry for each captured snapshot, then one or more snapshot entries in the ledger may be empty—e.g., no quarantine-related metadata may be recorded for one or more entries that have not been subject to a quarantine operation. The ledger **1500** may be configured to provide a source of truth for quarantine related information. That is, other data structures in the data management system that indicate and/or use quarantine-related metadata for snapshots may consult the ledger **1500** to determine a quarantine status for a snapshot.

FIG. **16** depicts an example recovery timeline that supports recording quarantining and release operations according to an example embodiment.

Recovery timeline **1600** depicts possible time-points for restoring a target object across a time period. The recovery timeline **1600** may include an indicator of possible restore points including snapshots (e.g., the snapshots **1615**) and point-in-time restore points, which may be referred to as PITs (e.g., the first PIT **1605-1** and the second PIT **1605-2**). Point-in-time restore points may enable a user to restore (e.g., in combination with a snapshot) a target object to a particular point-in-time with finer granularity than snapshots.

As described herein, creating snapshot-level restore points may involve capturing a state of a target object (e.g., a state of the file system, metadata, etc.), which may include placing the target object into a frozen state while the snapshot is generated. In some examples, snapshots are generated on a minute-level or hourly-level basis (e.g., down to 15 minute intervals).

Creating point-in-time level restore points may be less resource-intensive and less disruptive to the operation of a target object than creating snapshot-level restore points. In one example, continuous data protection techniques may create PITs on a per second or per millisecond basis (e.g., by capturing inputs/outputs for, e.g., data written to and read from, a memory disk) such that hundreds or thousands of PITs may be created between two snapshots. In such cases, a target object may be continuously restorable down to a second or millisecond, for example.

As described herein, a user may attempt to restore a target object (e.g., to a time that precedes a malware infection). In some examples, the user may indicate a particular point-in-time for recovery (e.g., a particular time-of-day to a millisecond), which may correspond to a particular PIT (e.g., the second PIT **1605-2**). As also described herein, quarantine information may be provided on a snapshot and/or file-level basis. Based on receiving a request to restore a target object to a particular point-in-time, the storage appliance may determine a quarantine status of a snapshot associated with the particular point-in-time—e.g., the snapshots that precede or follow the point-in-time. The storage appliance may determine the quarantine status based on consulting a ledger described herein, including with reference to FIG. **12** and FIG. **15**.

Based on determining the snapshot is quarantined, the storage appliance may identify a quarantine window that precedes and follows the snapshot. For example, the storage appliance may determine that the quarantine window extends from a preceding snapshot to a following snapshot—e.g., because it may be unknown when between the quarantined snapshot and the previous snapshot the malware infection occurred and when between the quarantined snap-

shot and a following snapshot (if not also quarantined) the malware infection was addressed. The PITs that fall within the quarantine window may be prevented from being used as restore points.

As an example, the storage appliance may receive a request from a user to restore a target object to the PIT **1605-1**. Based on receiving the request, the storage appliance may determine whether the first snapshot **1615-1** is quarantined (e.g., by consulting a ledger described herein, including with reference to FIG. 12 and FIG. 15). Additionally, or alternatively, the storage appliance may determine whether the second snapshot **1615-2** is quarantined. Based on determining that neither the first snapshot **1615-1** nor the second snapshot **1615-2** is quarantined, the storage appliance may determine that the first PIT **1605-1** is a safe recovery point and may restore the target object to the first point-in-time (e.g., using the first snapshot **1615-1** and a forward recovery procedure, using the second snapshot **1615-2** and a reverse recovery procedure).

In another example, the storage appliance may receive a request from a user to restore a target object to the second PIT **1605-2**. Based on receiving the request, the storage appliance may determine whether the second snapshot **1615-2** is quarantined. Additionally, or alternatively, the storage appliance may determine whether the third snapshot **1615-3** is quarantined. Based on determining that the third snapshot **1615-3** is quarantined, the storage appliance may generate the quarantine window **1610**, which extends from the second snapshot **1615-2** to the fourth snapshot **1615-4**. Based on determining that the second point-in-time occurs within the quarantine window **1610**, the storage appliance may prevent the target object from being restored to the second PIT **1605-2** (e.g., unless the user has administrative privileges for overriding the quarantine).

In some examples, a user may manually extend the quarantine window—e.g., by designating an unquarantined snapshot as quarantined based on suspecting or determining that the unquarantined snapshot is infected. For example, the user may extend the quarantine window **1610** by designating the second snapshot **1615-2** as quarantined such that the quarantine window **1610** may extend from the first snapshot **1615-1** to the fourth snapshot **1615-4**.

In some examples, instead of determining the quarantine window in response to a request to restore a PIT, the storage appliance may determine (e.g., periodically) the quarantine window throughout operation. In such cases, instead of first determining the quarantine window, the storage appliance may immediately determine whether a requested PIT is located within or outside of a quarantine window.

In some examples, a state of a PIT may be set based on determining whether a PIT is located within or outside of a quarantine window. For example, a state of PITs that occur outside of a quarantine window may be designated as recoverable and PITs that occur inside of the quarantine window may be designated as missed.

Although discussed in the context of PITs, quarantine window may be similarly generated and used to designate restore points within the quarantine window as unavailable for recovery. For example, on a snapshot-level basis, a quarantine window may be created between a first snapshot and an earlier snapshot such that any snapshots between the first and earlier snapshot may be prohibited for recovery.

The following set of examples describe various embodiments of methods, computer-readable media, and systems (e.g., machines, devices, or other apparatus) discussed herein.

A method is described. The method may include storing a plurality of snapshots of a target object that is within a computing system, the plurality of snapshots corresponding to a plurality of versions of the target object; storing an indication that information in one or more snapshots of the plurality of snapshots is flagged for quarantine; receiving a request to restore a snapshot of the target object, the snapshot of the target object included in the plurality of snapshots; determining, in response to the request to restore the snapshot of the target object, whether the snapshot of the target object is prohibited from being restored based at least in part on the indication that information in the one or more snapshots of the plurality of snapshots is flagged for quarantine; and determining whether to initiate a restoration operation in response to the request to restore the snapshot of the target object based at least in part on whether the snapshot of the target object is prohibited from being restored.

An apparatus is described. The apparatus may include a processor, memory coupled with the processor, and instructions stored in the memory. The instructions may be executable by the processor to cause the apparatus to store a plurality of snapshots of a target object that is within a computing system, the plurality of snapshots corresponding to a plurality of versions of the target object; store an indication that information in one or more snapshots of the plurality of snapshots is flagged for quarantine; receiving a request to restore a snapshot of the target object, the snapshot of the target object included in the plurality of snapshots; determine, in response to the request to restore the snapshot of the target object, whether the snapshot of the target object is prohibited from being restored based at least in part on the indication that information in the one or more snapshots of the plurality of snapshots is flagged for quarantine; and determine whether to initiate a restoration operation in response to the request to restore the snapshot of the target object based at least in part on whether the snapshot of the target object is prohibited from being restored.

Another apparatus is described. The apparatus may include means for storing a plurality of snapshots of a target object that is within a computing system, the plurality of snapshots corresponding to a plurality of versions of the target object; storing an indication that information in one or more snapshots of the plurality of snapshots is flagged for quarantine; receiving a request to restore a snapshot of the target object, the snapshot of the target object included in the plurality of snapshots; determining, in response to the request to restore the snapshot of the target object, whether the snapshot of the target object is prohibited from being restored based at least in part on the indication that information in the one or more snapshots of the plurality of snapshots is flagged for quarantine; and determining whether to initiate a restoration operation in response to the request to restore the snapshot of the target object based at least in part on whether the snapshot of the target object is prohibited from being restored.

A non-transitory computer-readable medium storing code is described. The code may include instructions executable by a processor to store a plurality of snapshots of a target object that is within a computing system, the plurality of snapshots corresponding to a plurality of versions of the target object; store an indication that information in one or more snapshots of the plurality of snapshots is flagged for quarantine; receiving a request to restore a snapshot of the target object, the snapshot of the target object included in the plurality of snapshots; determine, in response to the request to restore the snapshot of the target object, whether the

snapshot of the target object is prohibited from being restored based at least in part on the indication that information in the one or more snapshots of the plurality of snapshots is flagged for quarantine; and determine whether to initiate a restoration operation in response to the request to restore the snapshot of the target object based at least in part on whether the snapshot of the target object is prohibited from being restored.

In some examples of the method, apparatuses, and non-transitory computer-readable medium described herein, determining whether the snapshot of the target object is prohibited from being restored comprises: identifying, based at least in part on the request to restore the snapshot of the target object, a chain of snapshots for restoring the snapshot of the target object, the chain of snapshots including the snapshot of the target object; and determining whether information in one or more snapshots of the chain of snapshots has been flagged for quarantine.

In some examples of the method, apparatuses, and non-transitory computer-readable medium described herein, determining whether information in the one or more snapshots of the chain of snapshots has been flagged for quarantine comprises determining that there is no stored indication that any snapshots of the chain of snapshots is flagged for quarantine; and determining whether to initiate the restoration operation comprises initiating the restoration operation for the snapshot of the target object based at least in part on determining that there is no stored indication that any of the snapshots of the chain of snapshots is flagged for quarantine.

In some examples of the method, apparatuses, and non-transitory computer-readable medium described herein, determining whether information in the one or more snapshots of the chain of snapshots has been flagged for quarantine comprises obtaining, based at least in part on storing the indication, a first indication that information in a snapshot of the chain of snapshots is flagged for quarantine; and determining whether to initiate the restoration operation comprises preventing the restoration operation for the snapshot of the target object from being initiated based at least in part on obtaining the first indication.

Some examples of the method, apparatuses, and non-transitory computer-readable medium described herein may further include operations, features, means, or instructions for displaying an indication that the snapshot of the target object is prohibited from being restored based at least in part on preventing the restoration operation for the snapshot of the target object.

Some examples of the method, apparatuses, and non-transitory computer-readable medium described herein may further include operations, features, means, or instructions for receiving, in response to displaying the indication that the snapshot of the target object is prohibited, a command to restore the snapshot of the target object; determining that the command is associated with a user having privileges that satisfy a threshold; and restoring the snapshot of the target object into a sandboxed environment.

Some examples of the method, apparatuses, and non-transitory computer-readable medium described herein may further include operations, features, means, or instructions for initiating a partial restoration operation for the snapshot of the target object based at least in part on preventing the restoration operation for the snapshot of the target object, wherein the partial restoration operation prevents the information identified in the snapshot of the chain of snapshots from being restored.

In some examples of the method, apparatuses, and non-transitory computer-readable medium described herein, determining whether information in the one or more snapshots of the chain of snapshots has been flagged for quarantine comprises: determining whether a table comprises one or more entries identifying any snapshots of the chain of snapshots; and for snapshots identified by a plurality of entries, determining whether a most recent entry indicates a quarantining or a release from quarantine.

In some examples of the method, apparatuses, and non-transitory computer-readable medium described herein, determining whether information in the one or more snapshots of the chain of snapshots has been flagged for quarantine further comprises: determining that information in a snapshot of the chain of snapshots has been flagged for quarantine based at least in part on identifying an entry in the table that identifies that the snapshot of the chain of snapshots is currently flagged for quarantine.

In some examples of the method, apparatuses, and non-transitory computer-readable medium described herein, determining whether information in the one or more snapshots of the chain of snapshots has been flagged for quarantine comprises: determining whether a table comprises one or more entries identifying information included in any snapshots of the chain of snapshots.

In some examples of the method, apparatuses, and non-transitory computer-readable medium described herein, determining whether information in the one or more snapshots of the chain of snapshots has been flagged for quarantine comprises: determining whether a table comprises one or more entries identifying any snapshots of the chain of snapshots.

In some examples of the method, apparatuses, and non-transitory computer-readable medium described herein, determining whether the snapshot of the target object is prohibited from being restored comprises determining whether a table comprises an entry identifying the snapshot of the target object.

In some examples of the method, apparatuses, and non-transitory computer-readable medium described herein, determining whether the snapshot of the target object is prohibited from being restored comprises determining that information in the snapshot of the target object is not flagged for quarantine based at least in part on the stored indications; determining whether to initiate the restoration operation comprises determining to initiate the restoration operation for the snapshot of the target object based at least in part on determining that information in the requested snapshot of the target object is not flagged for quarantine; and the method further comprises initiating the restoration operation for the snapshot of the target object based at least in part on determining to initiate the restoration operation.

In some examples of the method, apparatuses, and non-transitory computer-readable medium described herein, determining whether the snapshot of the target object is prohibited from being restored comprises determining that information in the snapshot of the target object has been flagged for quarantine based at least in part on an indication of the stored indications indicating that the requested snapshot of the target object is flagged for quarantine; determining whether to initiate the restoration operation comprises determining not to initiate the restoration operation for the snapshot of the target object based at least in part on determining that information in the requested snapshot of the target object has been flagged for quarantine; and the method further comprises preventing the restoration operation

45

tion for the snapshot of the target object based at least in part on determining not to initiate the restoration operation.

Some examples of the method, apparatuses, and non-transitory computer-readable medium described herein may further include operations, features, means, or instructions for displaying an indication that the snapshot of the target object is prohibited from being restored based at least in part on preventing the restoration operation for the snapshot of the target object.

The terms “machine-readable medium,” “computer-readable medium” and “device-readable medium” mean the same thing and may be used interchangeably in this disclosure. The terms are defined to include both machine-storage media and transmission media. Thus, the terms include both storage devices/media and carrier waves/modulated data signals.

Although examples have been described with reference to specific example embodiments or methods, it will be evident that various modifications and changes may be made to these embodiments without departing from the broader scope of the embodiments. Accordingly, the specification and drawings are to be regarded in an illustrative rather than a restrictive sense. The accompanying drawings that form a part hereof, show by way of illustration, and not of limitation, specific embodiments in which the subject matter may be practiced. The embodiments illustrated are described in sufficient detail to enable those skilled in the art to practice the teachings disclosed herein. Other embodiments may be utilized and derived therefrom, such that structural and logical substitutions and changes may be made without departing from the scope of this disclosure. This detailed description, therefore, is not to be taken in a limiting sense, and the scope of various embodiments is defined only by the appended claims, along with the full range of equivalents to which such claims are entitled.

Such embodiments of the inventive subject matter may be referred to herein, individually and/or collectively, by the term “invention” merely for convenience and without intending to voluntarily limit the scope of this application to any single invention or inventive concept if more than one is in fact disclosed. Thus, although specific embodiments have been illustrated and described herein, it should be appreciated that any arrangement calculated to achieve the same purpose may be substituted for the specific embodiments shown. This disclosure is intended to cover any and all adaptations or variations of various embodiments. Combinations of the above embodiments, and other embodiments not specifically described herein, will be apparent to those of skill in the art upon reviewing the above description.

What is claimed is:

1. A method, comprising:

storing a plurality of snapshots of a target object that is within a computing system, the plurality of snapshots corresponding to a plurality of versions of the target object;

obtaining an indication that information associated with the target object is flagged for quarantine;

storing, based at least in part on the indication that the information associated with the target object is flagged for quarantine, one or more indications that the information, as included in one or more snapshots from among the plurality of snapshots, is flagged for quarantine;

46

receiving, after storing the one or more indications, a request to restore a snapshot of the target object, the snapshot of the target object included in the plurality of snapshots;

determining, in response to the request to restore the snapshot of the target object, whether the snapshot of the target object is prohibited from being restored based at least in part on the one or more stored indications that the information in the one or more snapshots from among the plurality of snapshots is flagged for quarantine; and

determining whether to initiate a restoration operation in response to the request to restore the snapshot of the target object based at least in part on whether the snapshot of the target object is prohibited from being restored.

2. The method of claim 1, wherein determining whether the snapshot of the target object is prohibited from being restored comprises:

identifying, based at least in part on the request to restore the snapshot of the target object, a chain of snapshots for restoring the snapshot of the target object, the chain of snapshots including the snapshot of the target object; and

determining whether information in one or more snapshots of the chain of snapshots has been flagged for quarantine.

3. The method of claim 2, wherein:

determining whether information in the one or more snapshots of the chain of snapshots has been flagged for quarantine comprises determining that there is no stored indication that any snapshots of the chain of snapshots is flagged for quarantine; and

determining whether to initiate the restoration operation comprises initiating the restoration operation for the snapshot of the target object based at least in part on determining that there is no stored indication that any of the snapshots of the chain of snapshots is flagged for quarantine.

4. The method of claim 2, wherein:

determining whether information in the one or more snapshots of the chain of snapshots has been flagged for quarantine comprises obtaining, based at least in part on reading at least one of the one or more stored indications, a first indication that information in a snapshot of the chain of snapshots is flagged for quarantine; and

determining whether to initiate the restoration operation comprises preventing the restoration operation for the snapshot of the target object from being initiated based at least in part on obtaining the first indication.

5. The method of claim 4, further comprising:

displaying an indication that the snapshot of the target object is prohibited from being restored based at least in part on preventing the restoration operation for the snapshot of the target object.

6. The method of claim 5, further comprising:

receiving, in response to displaying the indication that the snapshot of the target object is prohibited, a command to restore the snapshot of the target object; determining that the command is associated with a user having privileges that satisfy a threshold; and restoring the snapshot of the target object into a sandboxed environment.

7. The method of claim 4, further comprising:

initiating a partial restoration operation for the snapshot of the target object based at least in part on preventing the

47

restoration operation for the snapshot of the target object, wherein the partial restoration operation prevents the information identified in the snapshot of the chain of snapshots from being restored.

8. The method of claim 2, wherein determining whether information in the one or more snapshots of the chain of snapshots has been flagged for quarantine comprises:

determining whether a table comprises one or more entries identifying any snapshots of the chain of snapshots; and

for snapshots identified by a plurality of entries, determining whether a most recent entry indicates a quarantining or a release from quarantine.

9. The method of claim 8, wherein determining whether information in the one or more snapshots of the chain of snapshots has been flagged for quarantine further comprises:

determining that information in a snapshot of the chain of snapshots has been flagged for quarantine based at least in part on identifying an entry in the table that identifies that the snapshot of the chain of snapshots is currently flagged for quarantine.

10. The method of claim 2, wherein determining whether information in the one or more snapshots of the chain of snapshots has been flagged for quarantine comprises:

determining whether a table comprises one or more entries identifying information included in any snapshots of the chain of snapshots.

11. The method of claim 2, wherein determining whether information in the one or more snapshots of the chain of snapshots has been flagged for quarantine comprises:

determining whether a table comprises one or more entries identifying any snapshots of the chain of snapshots.

12. The method of claim 1, wherein:

determining whether the snapshot of the target object is prohibited from being restored comprises determining whether a table comprises an entry identifying the snapshot of the target object.

13. The method of claim 1, wherein:

determining whether the snapshot of the target object is prohibited from being restored comprises determining that information in the snapshot of the target object is not flagged for quarantine based at least in part on the one or more stored indications;

determining whether to initiate the restoration operation comprises determining to initiate the restoration operation for the snapshot of the target object based at least in part on determining that information in the snapshot of the target object is not flagged for quarantine; and the method further comprises initiating the restoration operation for the snapshot of the target object based at least in part on determining to initiate the restoration operation.

14. The method of claim 1, wherein:

determining whether the snapshot of the target object is prohibited from being restored comprises determining that information in the snapshot of the target object has been flagged for quarantine based at least in part on the one or more stored indications indicating that the snapshot of the target object is flagged for quarantine; determining whether to initiate the restoration operation comprises determining not to initiate the restoration operation for the snapshot of the target object based at least in part on determining that information in the snapshot of the target object has been flagged for quarantine; and

48

the method further comprises preventing the restoration operation for the snapshot of the target object based at least in part on determining not to initiate the restoration operation.

15. The method of claim 14, further comprising:

displaying an indication that the snapshot of the target object is prohibited from being restored based at least in part on preventing the restoration operation for the snapshot of the target object.

16. An apparatus, comprising:

a processor; and

a memory storing instructions that, when executed by the processor, cause the apparatus to:

store a plurality of snapshots of a target object that is within a computing system, the plurality of snapshots corresponding to a plurality of versions of the target object;

obtain an indication that information associated with the target object is flagged for quarantine;

store, based at least in part on the indication that the information associated with the target object is flagged for quarantine, one or more indications that the information, as included in one or more snapshots from among the plurality of snapshots, is flagged for quarantine;

receive, after storing the one or more indications, a request to restore a snapshot of the target object, the snapshot of the target object included in the plurality of snapshots;

determine, in response to the request to restore the snapshot of the target object, whether the snapshot of the target object is prohibited from being restored based at least in part on the one or more stored indications that the information in the one or more snapshots from among the plurality of snapshots is flagged for quarantine; and

determine whether to initiate a restoration operation in response to the request to restore the snapshot of the target object based at least in part on whether the snapshot of the target object is prohibited from being restored.

17. The apparatus of claim 16, wherein, to determine whether the snapshot of the target object is prohibited from being restored, the instructions are further executable by the processor to cause the apparatus to:

identify, based at least in part on the request to restore the snapshot of the target object, a chain of snapshots for restoring the snapshot of the target object, the chain of snapshots including the snapshot of the target object; and

determine whether information in one or more snapshots of the chain of snapshots has been flagged for quarantine.

18. The apparatus of claim 16, wherein, to determine whether the snapshot of the target object is prohibited from being restored, the instructions are further executable by the processor to cause the apparatus to:

determine whether a table comprises an entry identifying the snapshot of the target object.

19. The apparatus of claim 16, wherein:

to determine whether the snapshot of the target object is prohibited from being restored, the instructions are further executable by the processor to cause the apparatus to determine that information in the snapshot of the target object is not flagged for quarantine based at least in part on the one or more stored indications;

49

to determine whether to initiate the restoration operation, the instructions are further executable by the processor to cause the apparatus to determine to initiate the restoration operation for the snapshot of the target object based at least in part on determining that information in the snapshot of the target object is not flagged for quarantine; and

the instructions are further executable by the processor to cause the apparatus to initiate the restoration operation for the snapshot of the target object based at least in part on determining to initiate the restoration operation.

20. A non-transitory, computer-readable medium storing code comprising instructions executable by a processor of a device to cause the device to:

store a plurality of snapshots of a target object that is within a computing system, the plurality of snapshots corresponding to a plurality of versions of the target object;

obtaining an indication that information associated with the target object is flagged for quarantine;

store, based at least in part on the indication that the information associated with the target object is flagged

50

for quarantine, one or more indications that the information, as included in one or more snapshots from among the plurality of snapshots, is flagged for quarantine;

receive, after storing the one or more indications, a request to restore a snapshot of the target object, the snapshot of the target object included in the plurality of snapshots;

determine, in response to the request to restore the snapshot of the target object, whether the snapshot of the target object is prohibited from being restored based at least in part on the one or more stored indications that the information in the one or more snapshots from among the plurality of snapshots is flagged for quarantine; and

determine whether to initiate a restoration operation in response to the request to restore the snapshot of the target object based at least in part on whether the snapshot of the target object is prohibited from being restored.

* * * * *