

(12) **United States Patent**  
**Edwards et al.**

(10) **Patent No.:** **US 12,393,691 B2**  
(45) **Date of Patent:** **Aug. 19, 2025**

- (54) **SECURING NODE GROUPS**
- (71) Applicant: **Hewlett Packard Enterprise Development LP**, Spring, TX (US)
- (72) Inventors: **Nigel Edwards**, Bristol (GB); **Michael R. Krause**, Boulder Creek, CA (US); **Melvin Benedict**, Magnolia, TX (US); **Ludovic Emmanuel Paul Noel Jacquin**, Bristol (GB); **Luis Luciani**, Spring, TX (US); **Thomas Laffey**, Roseville, CA (US); **Theofrastos Koulouris**, Bristol (GB); **Shiva Dasari**, Austin, TX (US)
- (73) Assignee: **Hewlett Packard Enterprise Development LP**, Spring, TX (US)
- (\*) Notice: Subject to any disclaimer, the term of this patent is extended or adjusted under 35 U.S.C. 154(b) by 0 days.
- (21) Appl. No.: **18/528,893**
- (22) Filed: **Dec. 5, 2023**
- (65) **Prior Publication Data**  
US 2024/0104213 A1 Mar. 28, 2024
- Related U.S. Application Data**
- (63) Continuation of application No. 17/280,507, filed as application No. PCT/US2019/012610 on Jan. 8, 2019, now Pat. No. 11,868,474.
- (51) **Int. Cl.**  
**G06F 21/00** (2013.01)  
**G06F 21/32** (2013.01)  
(Continued)
- (52) **U.S. Cl.**  
CPC ..... **G06F 21/57** (2013.01); **G06F 21/32** (2013.01); **H04L 9/0816** (2013.01); **H04L 9/3226** (2013.01)

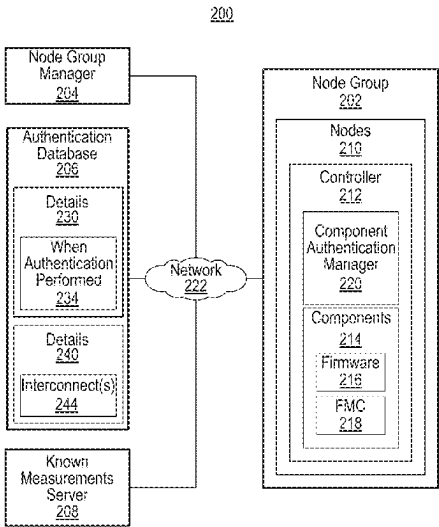
- (58) **Field of Classification Search**  
None  
See application file for complete search history.
- (56) **References Cited**  
**U.S. PATENT DOCUMENTS**  
8,738,894 B2 5/2014 Batke et al.  
9,923,918 B2 3/2018 Nicodemus et al.  
(Continued)  
**FOREIGN PATENT DOCUMENTS**  
CN 1868189 A 11/2006  
CN 102123031 A 7/2011  
(Continued)  
**OTHER PUBLICATIONS**  
Oliveira, Ana Cristina et al. Efficient network service level agreement monitoring for cloud computing systems. 2014 IEEE Symposium on Computers and Communications (ISCC). <https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=6912551> (Year: 2014).\*  
(Continued)

*Primary Examiner* — Jeremiah L Avery  
(74) *Attorney, Agent, or Firm* — Trop, Pruner & Hu, P.C.

(57) **ABSTRACT**

A method for securing a plurality of compute nodes includes authenticating a hardware architecture of each of a plurality of components of the compute nodes. The method also includes authenticating a firmware of each of the plurality of components. Further, the method includes generating an authentication database comprising a plurality of authentication descriptions that are based on the authenticated hardware architecture and the authenticated firmware. Additionally, a policy for securing a specified subset of the plurality of compute nodes is implemented by using the authentication database.

**20 Claims, 11 Drawing Sheets**



- (51) **Int. Cl.**  
**G06F 21/57** (2013.01)  
**H04L 9/08** (2006.01)  
**H04L 9/32** (2006.01)

(56) **References Cited**

U.S. PATENT DOCUMENTS

9,998,438	B2	6/2018	Kannan et al.	
10,484,418	B2 *	11/2019	Ahuja	H04L 63/105
10,542,044	B2 *	1/2020	Vissamsetty	H04L 63/1441
10,721,275	B2 *	7/2020	Kung	H04L 63/20
10,771,264	B2	9/2020	Edwards et al.	
11,750,412	B2 *	9/2023	Ansari	H04L 67/125
				709/225
2005/0278775	A1	12/2005	Ross	
2006/0200859	A1	9/2006	England et al.	
2011/0029769	A1	2/2011	Aissi et al.	
2011/0154501	A1	6/2011	Banginwar et al.	
2011/0185181	A1	7/2011	Lin	
2014/0109076	A1 *	4/2014	Boone	H04L 63/105
				717/170
2014/0365755	A1	12/2014	Shutt et al.	
2014/0380425	A1 *	12/2014	Lockett	G06F 21/00
				726/4
2015/0052596	A1 *	2/2015	Ayanam	H04W 4/38
				726/8
2015/0095648	A1 *	4/2015	Nix	H04L 12/2854
				713/170
2016/0371493	A1	12/2016	Prakash et al.	
2017/0010899	A1 *	1/2017	Dasar	G06F 8/654
2017/0195124	A1 *	7/2017	Obaidi	H04L 9/0861
2017/0201545	A1	7/2017	Nicodemus et al.	
2018/0075242	A1 *	3/2018	Khatri	H04L 9/3263
2018/0367541	A1 *	12/2018	Ponnuswamy	H04L 63/101
2019/0188389	A1 *	6/2019	Peled	H04L 63/20
2019/0373021	A1 *	12/2019	Parthasarathy	G06F 9/45558
2020/0004742	A1 *	1/2020	Nagarajan	G06F 16/248
2020/0007583	A1 *	1/2020	Dixit	H04L 63/102
2020/0036593	A1 *	1/2020	Sethi	H04L 41/0895
2020/0162503	A1 *	5/2020	Shurtleff	G06F 11/3013
2023/0106581	A1 *	4/2023	Sood	H04L 63/166
				709/203

FOREIGN PATENT DOCUMENTS

CN	105515776	A	4/2016
CN	108345805	A	7/2018
EP	2378454	A2	10/2011
WO	2013/028059	A1	2/2013

OTHER PUBLICATIONS

Aldinucci, Marco et al. Autonomic QoS in ASSIST grid-aware components. 14th Euromicro International Conference on Parallel, Distributed, and Network-Based Processing (PDP'06). [https://](https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=1613276)

[ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=1613276](https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=1613276) (Year: 2006).\*

Hwang, JeeHyun et al. ACPT: A Tool for Modeling and Verifying Access Control Policies. 2010 IEEE International Symposium on Policies for Distributed Systems and Networks. <https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=5629938> (Year: 2010).\*

Jaidi, Faouzi; Ayachi, Faten Labbene. An Approach to Formally Validate and Verify the Compliance of Low Level Access Control Policies. 2014 IEEE 17th International Conference on Computational Science and Engineering. <https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=7023798> (Year: 2014).\*

Arthur et al., "A Practical Guide to TPM 2.0", ApressOpen, 2015, 375 pages.

B. Moran et al., "A Firmware Update Architecture for Internet of Things Devices," Mar. 2, 2018, 17 Pgs. Retrieved Aug. 10, 2018.

Cisco, "Cisco Trust Anchor Technologies," Data Sheet, 2015, 5 Pgs.

Dave Heller et al., "Using Trusted Boot on IBM OpenPOWER Servers," Feb. 17, 2017, 16 Pgs., Retrieved Aug. 10, 2018.

Google Cloud, "Titan in Depth: Security in Plaintext", available online at <<https://cloud.google.com/blog/products/identity-security/titan-in-depth-security-in-plaintext>>, Aug. 25, 2017, 6 pages.

IEEE Standard Association, "IEEE Standard for Local and Metropolitan Area Networks—", Secure Device Identity, IEEE Std 802.1AR™—, 2018, 73 pages.

Intel, "PCI Express Device Security Enhancements", version 0.7, Jun. 2018, 44 pages.

International Search Report and Written Opinion received for PCT Patent Application No. PCT/US19/12610, mailed on Oct. 7, 2019, 10 pages.

Kelly et al., "Project Cerberus Security Architecture Overview Specification", Open Compute Project, Microsoft Corporation, 2017, 17 pages.

Kurachi, Ryo et al. Improving secure coding rules for automotive software by using a vulnerability database. 2018 IEEE International Conference on Vehicular Electronics and Safety (ICVES). <https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=8519496> (Year: 2018).

Lorch, Markus et al. A hardware-secured credential repository for Grid PKIs. IEEE International Symposium on Cluster Computing and the Grid, 2004. CCGrid 2004. <https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=1336679> (Year: 2004).

Park, Jin Hyung et al. Security Architecture for a Secure Database on Android. IEEE Access, vol. 6. <https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=8271994> (Year: 2018).

Security Boulevard, "Implicit Identity Based Device Attestation", TCG Admin, available online at <<https://securityboulevard.com/2018/03/implicit-identity-based-device-attestation/>>, Mar. 5, 2018, 6 pages.

TCG, "TCG Platform Attribute Credential Profile", Specification Version 1.0, Revision 16, Jan. 16, 2018, 43 pages.

TCG, "TPM Keys for Platform DevID for TPM2", Specification Version 0.7, Revision 45, Dec. 12, 2018, 62 pages.

USB Type-C Authentication Specification, "Universal Serial Bus Type-C™ Authentication Specification", Revision 1.0 with ECN and Errata, Jul. 24, 2017, pp. 1-66.

\* cited by examiner

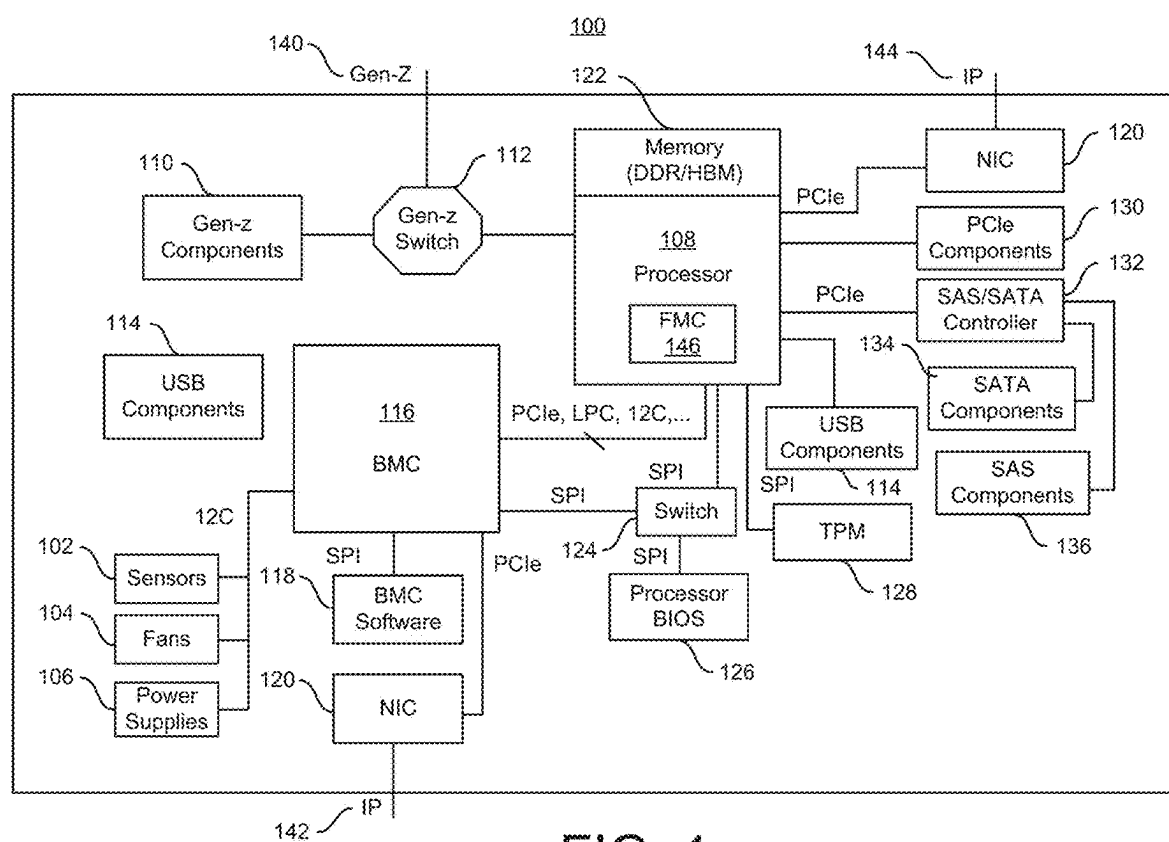


FIG. 1

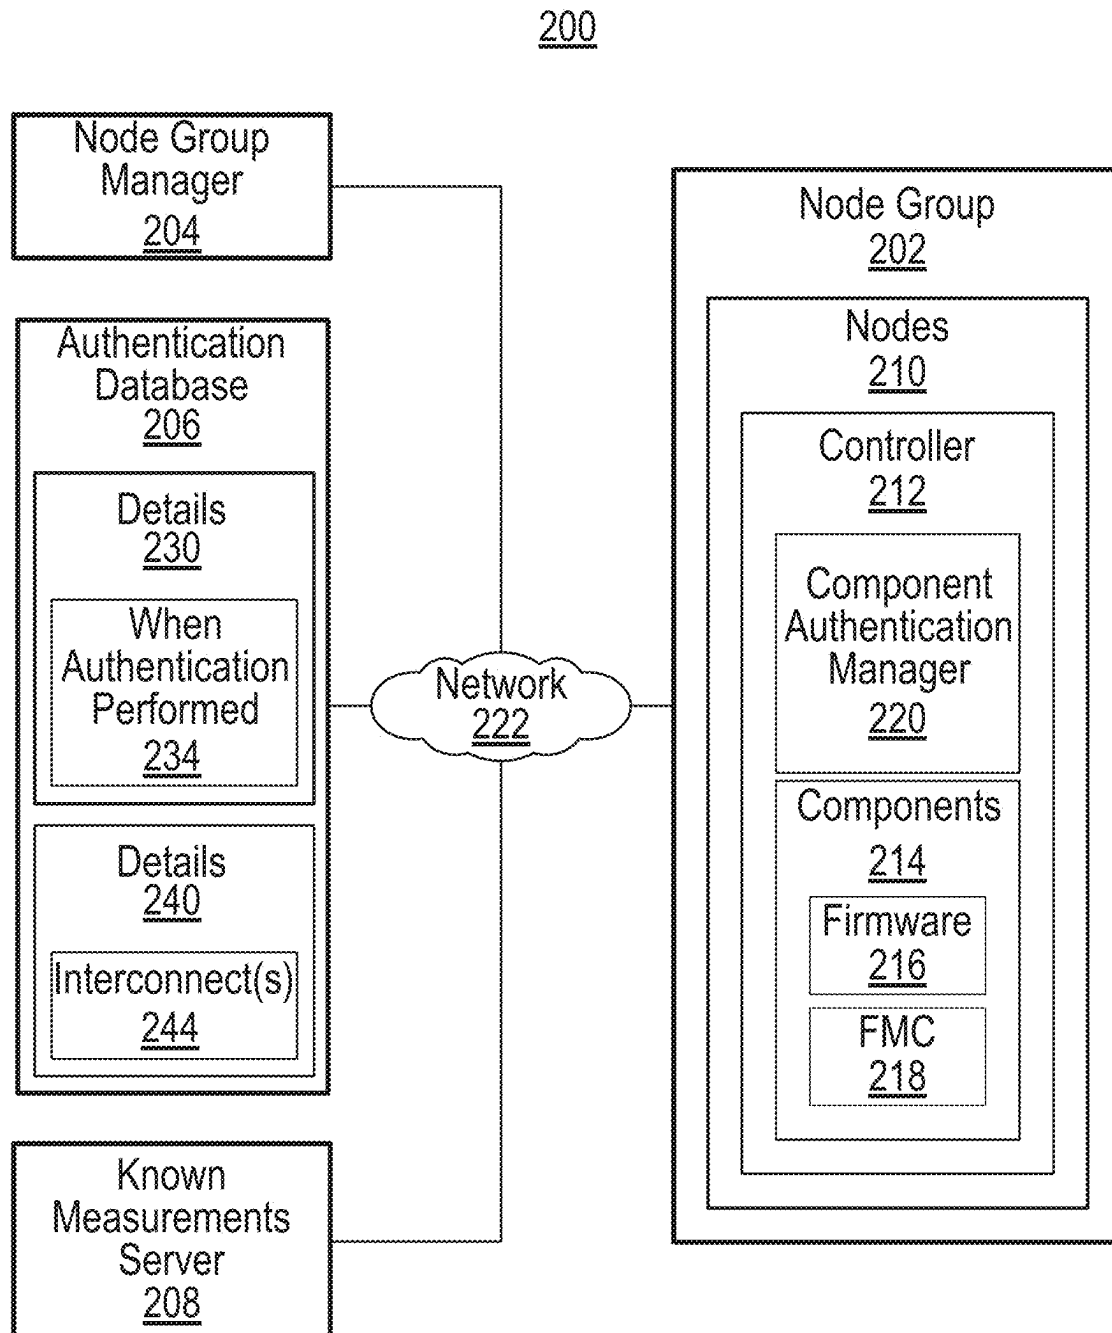


FIG. 2




FIG. 3

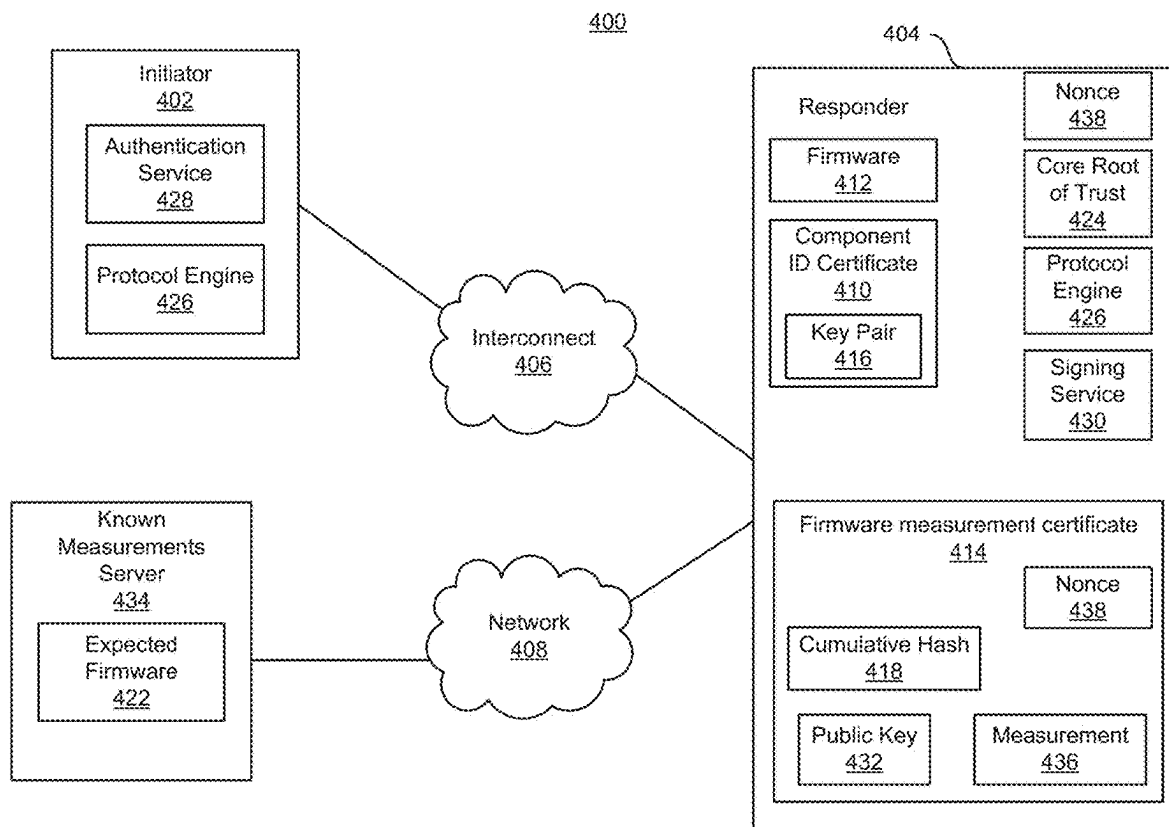
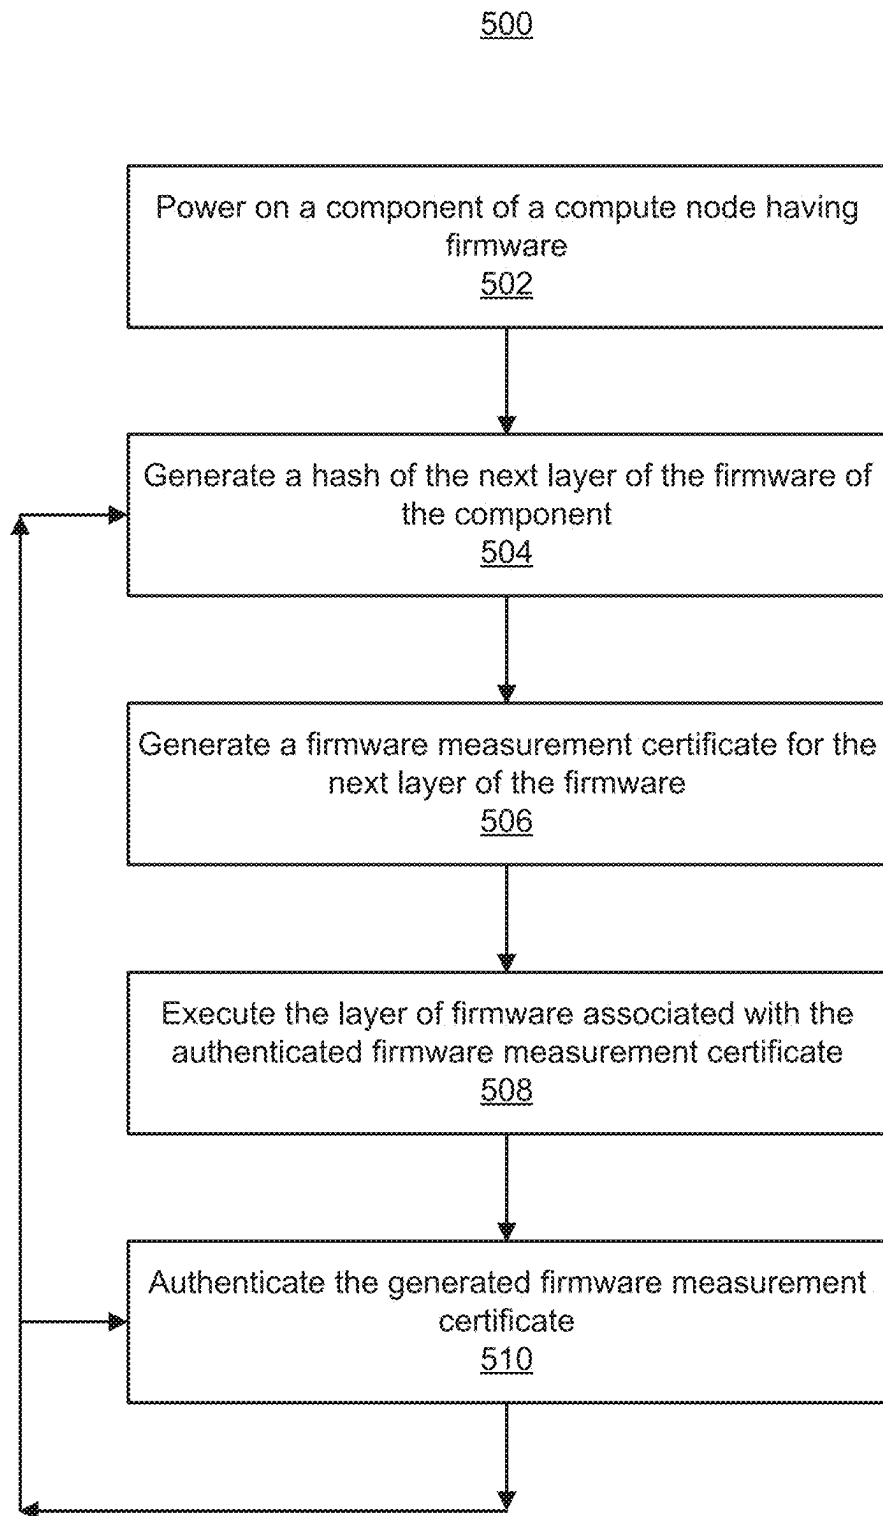


FIG. 4

**FIG. 5**

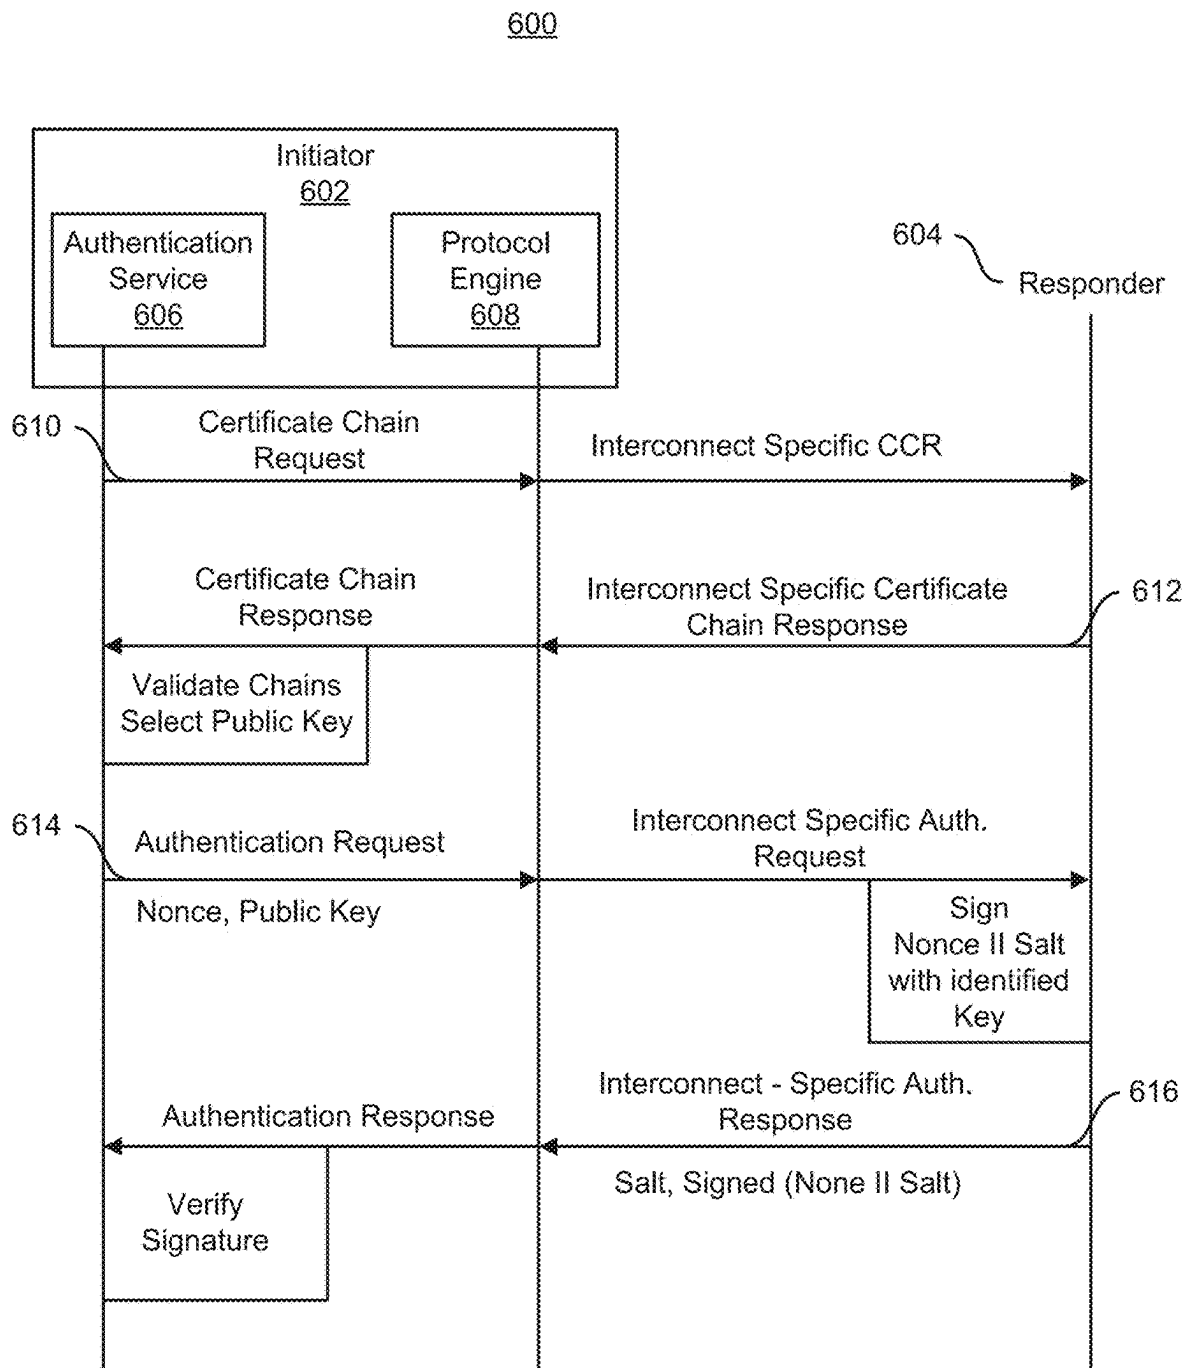


FIG. 6



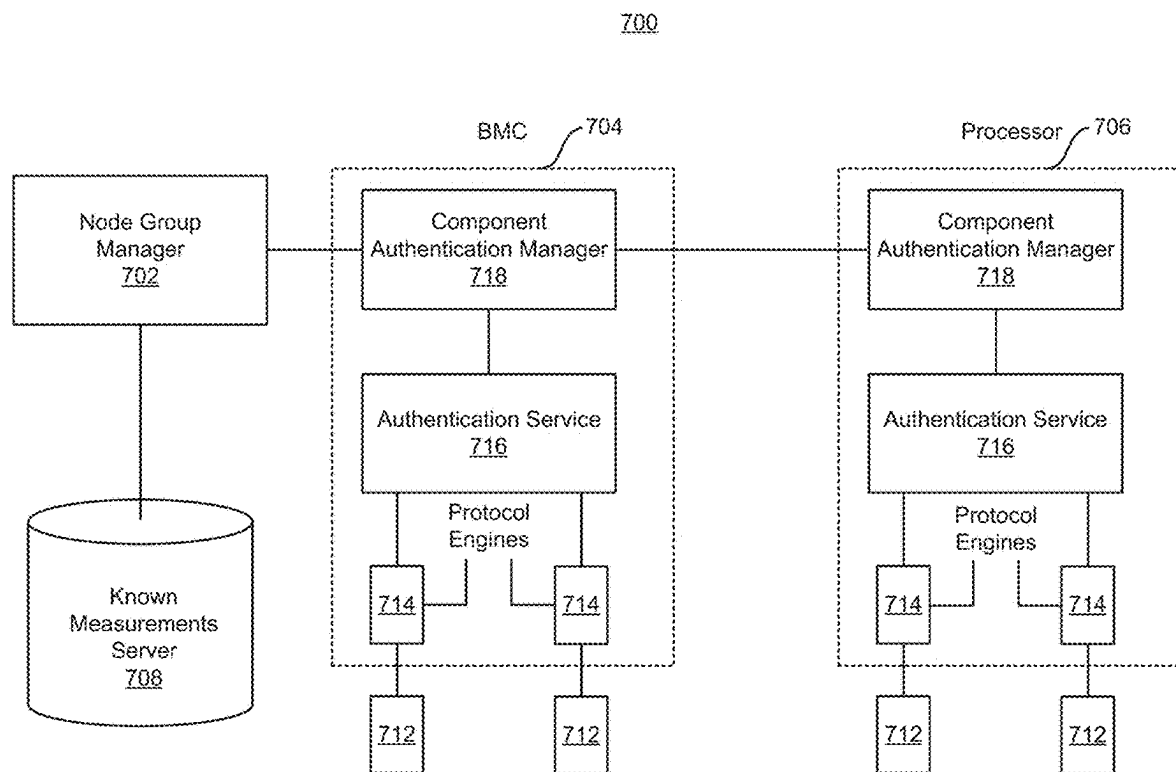


FIG. 7

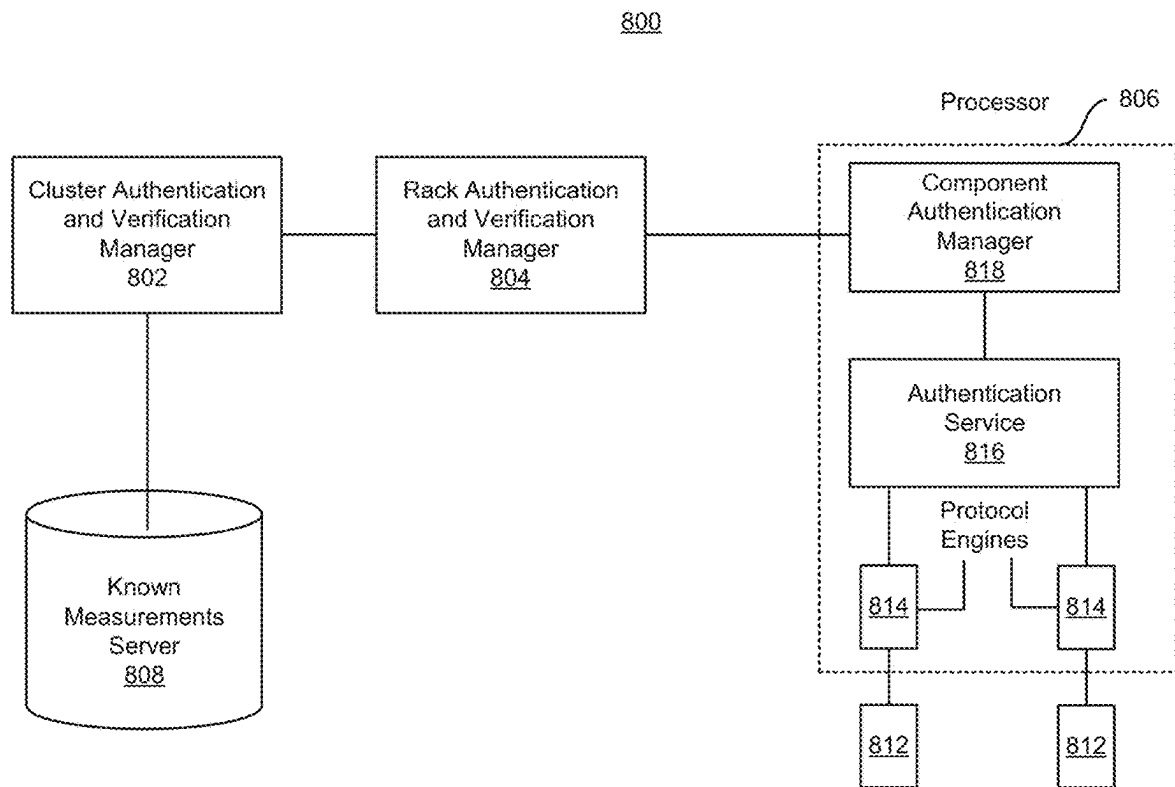


FIG. 8

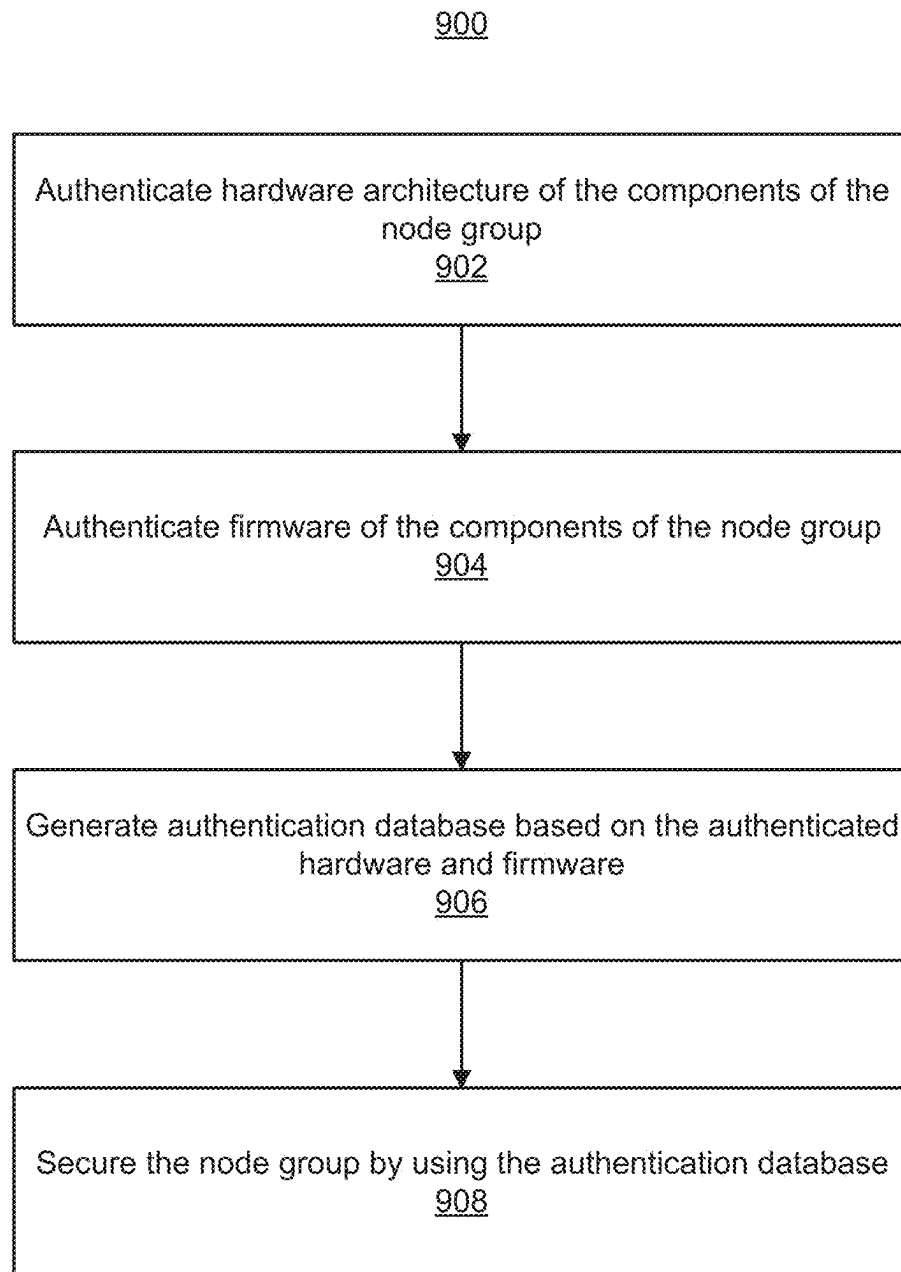


FIG. 9

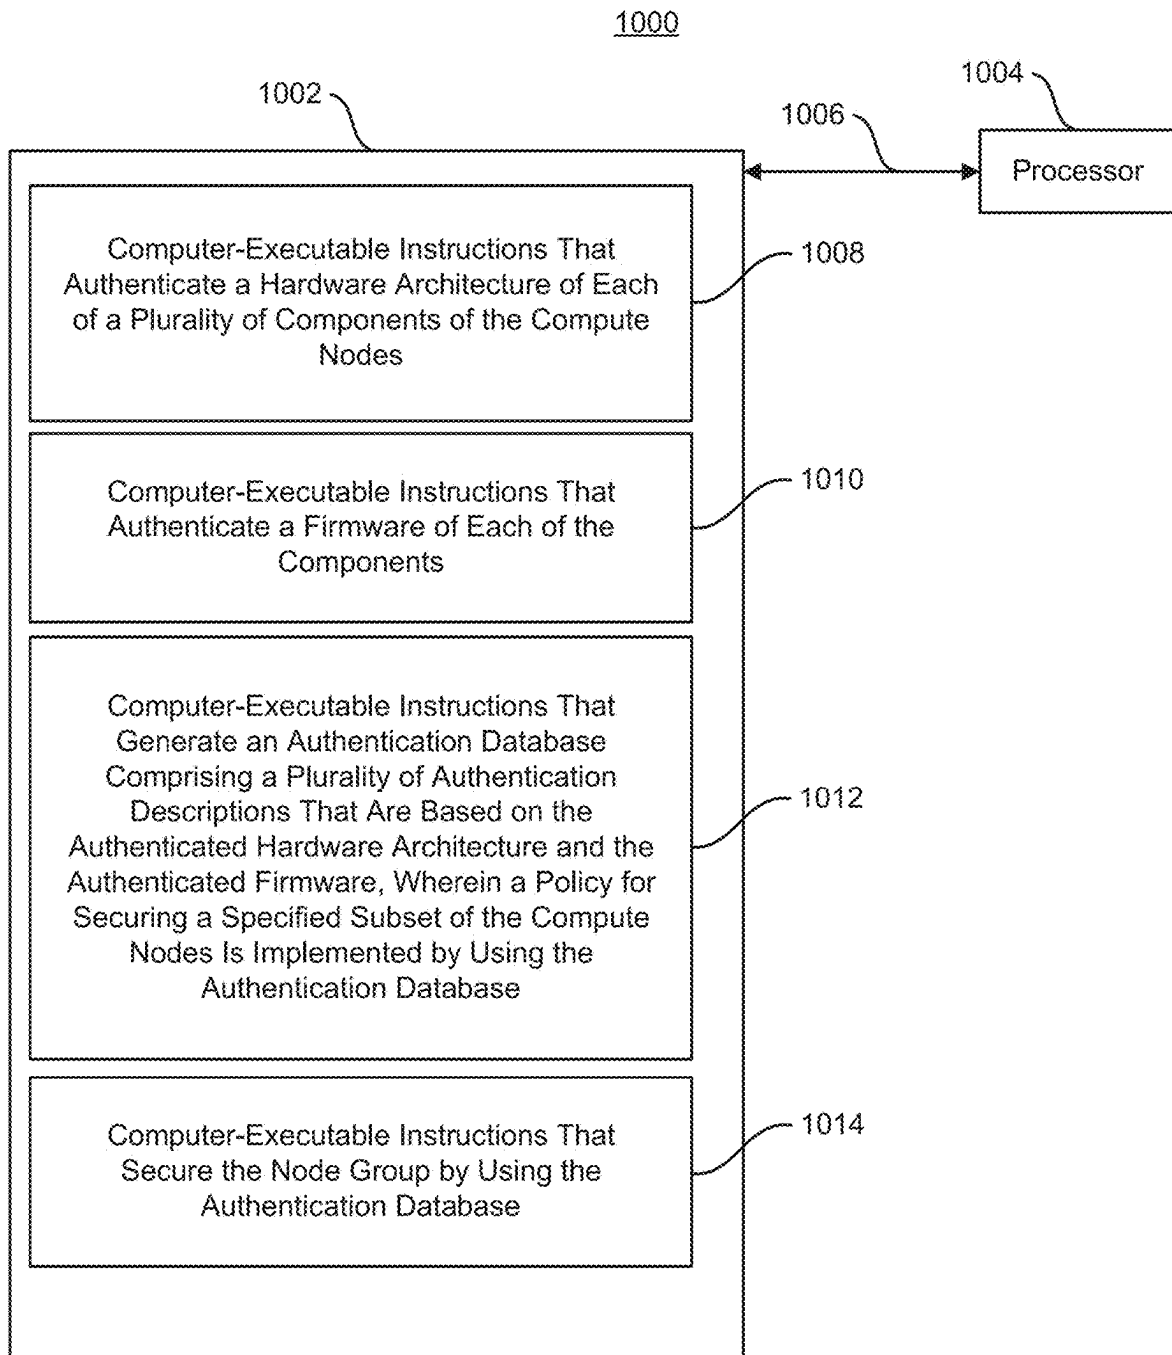


FIG. 10

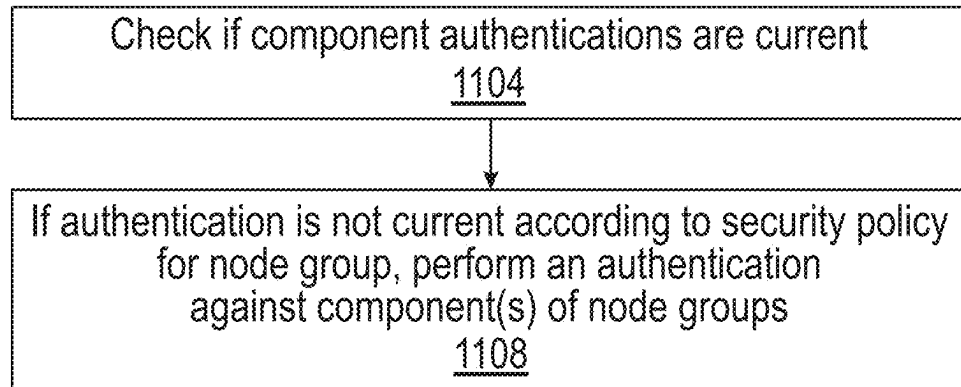


FIG. 11

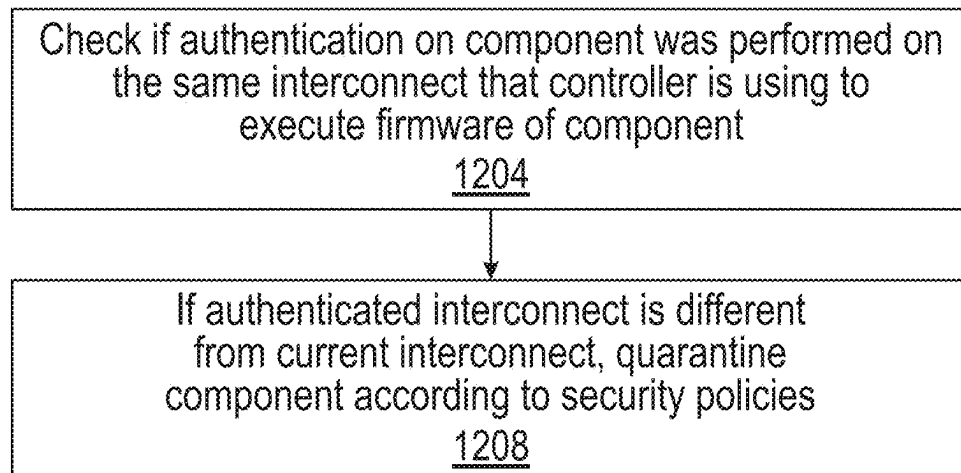


FIG. 12

1

**SECURING NODE GROUPS****CROSS-REFERENCE TO RELATED APPLICATIONS**

This application claims priority to U.S. patent application Ser. No. 17/280,507, which was filed on Mar. 26, 2021, which is based on Application No. PCT/US2019/012610 filed on Jan. 8, 2018, which is herein included by reference in its entirety for all purposes.

**BACKGROUND**

Within a chassis or enclosure of a compute node, a computer system, or host, there may be hundreds of plug-  
 5 gable components, from temperature sensors and power supplies to memory modules and processors. Within a rack or a cluster of compute nodes, there may be thousands of such components. However, each component may represent a security vulnerability, i.e., a potential attack vector. A component may be a potential attack vector if the compo-  
 10 nent is counterfeit or contains malware that may compromise the compute node. One potential way to compromise a component is to corrupt the firmware that is used to operate the component. Even simple components, such as fans and sensors, if compromised, can cause damage to a compute node through overheating or fire. Hence, identifying com-  
 15 promised components may be useful for preventing the components' misuse.

**BRIEF DESCRIPTION OF THE DRAWINGS**

The present disclosure may be understood from the fol-  
 20 lowing detailed description when read with the accompanying Figures. In accordance with the standard practice in the industry, various features are not drawn to scale. In fact, the dimensions of the various features may be arbitrarily increased or reduced for clarity of discussion.

Some examples of the present application are described with respect to the following figures.

FIG. 1 is an example compute node of a node group that may be secured.

FIG. 2 is an example system for securing node groups.

FIG. 3 is an example firmware measurement certificate for securing node groups.

FIG. 4 is an example system for securing node groups.

FIG. 5 is a process flow diagram of a method for generating firmware measurement certificates to secure node groups.

FIG. 6 is a message flow diagram for authenticating components in a system for securing node groups.

FIG. 7 is an example system for securing node groups.

FIG. 8 is an example system for securing node groups.

FIG. 9 is a process flow diagram of a method for securing node groups.

FIG. 10 is an example system comprising a tangible, non-transitory computer-readable medium that stores code for securing node groups.

FIGS. 11 and 12 are flow diagrams associated with a node group manager.

**DETAILED DESCRIPTION**

Examples of authentication include universal serial bus (USB) Type C authentication, which enables a compute node, i.e., a host, to authenticate compliant USB compo-  
 65 nents. USB Type C authentication also forms the basis for

2

potential peripheral component interconnect express (PCIe) authentication mechanisms, which allow PCIe components to be authenticated. The pattern of authentication in USB Type C and PCIe can be extended to internal buses, and other protocols and interconnects.

The purpose of component authentication is to establish trust in the component. The authentication mechanisms discussed above may establish that the hardware of the component is from a known and trusted manufacturer.  
 10 However, establishing that a component is from a known and trusted manufacturer does not mean that the firmware running inside the component is correct and trustworthy. Being correct may mean that the correct firmware and the correct version of the firmware are installed in the compo-  
 15 nent. Being trustworthy may mean that the firmware can be trusted not to breach the security of the component on which the firmware is running.

Additionally, while the authentication of each component may be challenging, the challenge may be greater when considering that multiple components may be included in one compute node. Further, compute nodes may be combined into node groups, such as a chassis enclosure with two compute nodes, or a rack system with all the compute nodes in a rack of blade servers. Node groups may also include node clusters, which may include hundreds of compute nodes or more, which may include tens of thousands of components.

However, such a scale may introduce additional chal-  
 30 lenges. For example, at the node cluster scale, the collective rate at which components fail and are replaced may be such that the system configurations of many compute nodes, even of the same type and from the same manufacturer, may vary from their original factory setups. Accordingly, attempting to use automated methods to authenticate the component and verify the firmware of the components in a cluster may be challenging. The varying configurations may reduce assumptions that may be made about the original factory setups at the cluster scale. Thus, automated methods, which may be more efficient if such assumptions could be made at the cluster scale may not be useful. Rather, authentication and verification may become customized, which may be tedious and costly.

Further, in a node cluster, different compute nodes with the same hardware components may run different versions of firmware at any given time. For example, as a practical consideration, firmware updates for the components may be carried out in a staged fashion. This may ensure the availability of the node cluster even when some components are unavailable due to firmware updates. Thus, different sets of components may have different versions of firmware until all the stages of a firmware update in the node cluster is complete. Additionally, the same types of components from the same manufacturer may arrive with different versions of factory-installed firmware, whether initially installed or replaced. Additionally, a node cluster is a complex computer system, whereby higher-level data center and node cluster management systems may be tasked with creating logical pools of resources from the physical collections of compo-  
 50 nents present in a node cluster. As such, this may lead to some parts of the node cluster being reconfigured, rebooted or taken offline more often than others. In some scenarios, rebooting and taking compute nodes offline may be opportunities when the components on a compute node are authenticated, or updated with new versions of firmware. Thus, if different compute nodes within a node cluster reboot or go offline at different times, the components on those

compute nodes may end up with different versions of firmware, some of which may be compromised.

Accordingly, in examples of the present disclosure, a node cluster-scale component authentication and verification system may be provided to dynamically manage such discrepancies at scale by providing the abilities to identify, report, and manage components on groups of nodes according to specified policies. In this way, the component authentication and verification system may identify components within node groups that may be running vulnerable or defective versions of firmware.

FIG. 1 is an example compute node **100** of a node group that may be secured. The node group may be authenticated by ensuring the components of each of the example compute nodes **100** in the node group are trustworthy. The example compute node **100** includes multiple components that may be authenticated along with the firmware installed on each of the components. Firmware may be computer instructions that operates the various components of the compute node **100**. In examples, some components may be installed by a manufacturer of the compute node **100**. Additionally, some components may be field replaceable, meaning the components are replaced after purchase of the compute node **100** by installing the components on the compute node **100** when the compute node **100** is powered off. Further, some components may be hot-plugged. Being hot-plugged means that the component is physically connected to an interconnect of the compute node **100** while the compute node **100** is powered on. The compute node **100** may include components with a range of capabilities, including components with little to no processing ability, such as sensors **102**, fans **104**, and power supplies **106**, and components with complex processing capabilities, such as a general-purpose processor **108**. Additional components of the compute node **100** may include, for example, Gen-Z components **110**, a Gen-Z switch **112**, USB components **114**, a baseboard management controller (BMC) **116**, BMC software **118**, multiple network interface controllers (NICs) **120**, memory **122**, a serial peripheral interconnect (SPI) switch **124**. The SPI switch **124** may be in the bus that is used to access read-only memory (ROM). The SPI switch **124** may enable the BMC **116** to check if the general-purpose processor **108** is loading the correct firmware. For example, The SPI switch **124** may enable the BMC **116** to read the basic input output system (BIOS) **126**, before the BIOS **126** is loadable by the processor **108**. If the BIOS **126** is correct, the BMC **116** may flip the switch to allow the processor **108** to load the BIOS **126**. The SPI switch **124** may also enable the BMC **116** to restore or update BIOS **126**. Further, the components may include the BIOS **126**, a trusted platform module **128**, PCIe components **130**, a serial attached SCSI (SAS)/SATA controller **132**, SATA components **134**, and SAS components **136**. These components may also be connected over a range of interconnects, including inter-integrated circuit (I2C), PCIe, USB, double data rate (DDR), high bandwidth memory (HBM), and Gen-Z, within the compute node **100**.

The example compute node **100** may be a network switch, router, network server, and the like. As such, the components of the example compute node **100** may vary, including fewer or additional components. For example, the compute node **100** may include multiple processors **108**, or a platform controller hub. Additionally, a compute node **100** may include one or more processors **108** but not the BMC **116**. Accordingly, authentication and verification may be conducted out-of-band via the BMC **116**. Alternatively, authentication and verification may be conducted in-band via the processor **108**.

In examples of the present disclosure, one or more of the components of the system **100** may include firmware measurement certificates **146**, which may be used to authenticate that the firmware for the components of the compute node are correct and trustworthy. The firmware measurement certificates **146** may provide trustworthy measurements of the firmware running on the components. A measurement is the value of binary image of the firmware that is loaded into memory for execution. Thus, an example authentication and verification system may use the firmware measurement certificates **146** to determine if the associated components are running firmware according to the manufacturers' specifications, or if the firmware may be compromised. The example compute node **100** may communicate with other compute nodes in the same node group, or on different networks, across Gen-Z or IP networks. For example, authentication and verification of components may be managed from outside the compute node **100**, on a rack, node cluster, or fabric management system. Authentication managed from outside of the compute node **100** may take place over a Gen-Z interconnect **140**, or IP networks **142**, **144**.

The hardware and firmware of the components may be authenticated in response to a request from an authentication initiator, also referred to herein as the initiator. The initiator may authenticate a computer component, e.g., the USB components **114**, by executing a series of calls to the component, also referred to herein as a responder. Examples of initiators may include software or firmware executing on the compute node **100**. Example initiators may include an operating system executing on the processors **108**, and the firmware of the BMC **116**.

FIG. 2 is an example system **200** for securing node groups. The example system **200** may include a node group **202**, a node group manager **204**, an authentication database **206**, and a known measurements server **208**. The node group **202** may be a chassis enclosure, a rack, a node cluster, or other grouping of compute nodes. The node group **202** may include multiple nodes **210**, which may be compute nodes, such as switches, routers, blade servers, and the like. The nodes **210** may include a controller **212** and one or more components **214**. The components **214** may include various compute node components, such as the components described with respect to FIG. 5. The components may include firmware **216** and one or more firmware measurement certificates (FMCs) **218**. The firmware **216** may be the instructions installed on the component **214** that operate the component **214**. The FMC **218** may be similar to a digital certificate, which is an electronic document that may be distributed by an issuing authority. Digital certificates may ensure the trustworthiness of a component **214**, a compute node **210**, and the like. The firmware measurement certificate **218** may enable the component authentication manager **220** to ensure the trustworthiness of the firmware **216** by providing an accurate measurement of the firmware **216** that is loaded in memory for operating the component **214**. In examples, the measurements in the firmware measurement certificates **218** may be compared to measurements in the known measurements server **208**. The known measurements server **208** may include measurements of the firmware provided by the component manufacturers, indicating the measurements of the binary image of the firmware that is installed on component **214** during manufacture, or during a legitimate update from the manufacturer. Thus, the comparison may be useful in determining whether the firmware **216** is trustworthy or compromised.

The controller **212** may be a BMC or a processor, such as the BMC **116** and processor **108**. Example nodes **210** may

include multiple controllers **212**, including combinations of BMCs and processors. The controller **212** includes a component authentication manager **220**, which may be firmware that performs authentications of the components **214** on all of the nodes **212** of the group when the controller **212** powers up, exits a low-power state, or on demand. The component authentication manager **220** may perform authentication of the hardware of the component **214**, and the firmware **216**. Thus, the controller **212** may represent an initiator, and the components may represent the responders. In examples, the controller **212** may authenticate itself and another controller **212**. For example, the BMC **116** may authenticate itself, another BMC **116**, and the processors **108**.

The component authentication manager **220** may also provide an API of operations that may be used to execute authentications of individual or groups of components **214**. Further, when performing authentications, the component authentication manager **220** may store details of the authentications in the authentication database **206**. Details **230** about the authentication may include, for example, as depicted at **234**, when an authentication is performed. Details **240** in an example node **210** may specify one or more interconnects **244**: PCIe, low pin count (LPC), Inter-Integrated Circuit (I2C), between the processor **108** and the BMC **116** over which the authentication is performed. Accordingly, the API may be used to query the authentication database **206** for details about the authentications.

The node group manager **204** may verify the authentications of the components **214** in node groups **202** by issuing calls to the API of the component authentication manager **220**. Referring also to FIG. **11** in conjunction with FIG. **2**, the node group manager **204** may issue these calls within scripts and computer programs that automatically enforce security policies for node groups **202** by checking (block **1104**), for example, if components' authentications are current. As depicted in block **1108**, if the authentication is not current according to a security policy for the node group **202**, the node group manager **204** may perform an authentication against one or more components **214**. Referring to FIG. **12** in conjunction with FIG. **2**, in another example, the node group manager **204** may use the component authentication manager **220** API to check (block **1204**) if the authentication on a specific component **214** was performed over the same interconnect that the controller **212** is currently using to execute the firmware **216** of the component **214**. If the authenticated interconnect is different from the current interconnect, the component **214** may be compromised. Accordingly, in such an example and as depicted in block **1208**, the node group manager **204** may take additional steps to quarantine the potentially compromised component **214** according to predetermined security policies.

The node group **202**, node group manager **204**, authentication database **206** and known measurements server **208** may be in communication over a network **222**. The network **222** may be a computer communication network or collection of networks, including a local area network, wide area network, the Internet, and the like.

FIG. **3** is an example firmware measurement certificate **300** in a system for securing node groups. The firmware measurement certificate **300** is an example of the firmware measurement certificate **218**, which may be used for securing the firmware **216** of each of the components **214**. Referring back to FIG. **3**, the firmware measurement certificate **300** includes an issuer **302** and attributes **304**. The issuer **302** may be the name of this issuer. This name may be associated with a public key that is signed by a certificate

issued to the issuer, with the root being self-signed. The root certification authority key may be known. There may be no higher authority than the root certification authority, which represents a trust anchor that is known to the authentication initiator. The trust anchor may be represented in a local data store that identifies trustworthy certification authorities. For example, when identity certificates are used for a web browser, the web browser provider may configure trust anchors into the web browser before releasing it for general use. Similarly, the initiator may have trust anchors, i.e., one or more known root certification authorities that the manufacturer of the initiator trusts. The signature may be applied to the whole certificate but is a separate structure (not shown). In examples, the issuer **302** may represent a core root of trust or a specific layer of firmware **216**. The attributes **304** may include a component ID **306**, a cumulative hash **308**, measurements **310**, an alias ID **312**, and a nonce **314**. Similar to the subject of an identity certificate, the alias ID **312** may be a public key that identifies the owner of the firmware measurement certificate **300**. The alias ID **312** may be authenticated by the initiator against the trust anchors. To limit the authentication of the firmware measurement certificate **300** to one layer of the firmware **216**, the alias ID **312** may be rendered unusable to higher layers of the firmware **216**. The component ID **306** identifies a public-private key pair assigned to the component **214**. In examples, the component ID **306** may be used to sign the first firmware measurement certificate in the hierarchy, i.e., layer **L0** of the firmware **216**. The firmware measurement certificates **300** for multiple layers of firmware **216** may form a hierarchy, wherein each firmware measurement certificate **300** is issued by the alias ID in the firmware measurement certificate **300** for the previous layer of firmware **216**. The component ID **306** may be used to sign the first firmware measurement certificate **300** in the hierarchy and subsequent firmware measurement certificates **300** may be signed by the alias ID **312** in the previous layer of the firmware **216**. In examples, the component ID **306** may not be accessible outside the core root of trust.

The cumulative hash **308** may be cryptographic hash representation of all layers of firmware **216** up to the layer of firmware **216** being secured. An example equation for calculating the cumulative hash **308** over layers **0** through **n** of the firmware **216** is shown in Equation 1. In Equation 1,  $H_{-}$  represents a cumulative hash function, and  $H$  represents a hash function that is approved by the National Institute of Standards and Technology (NIST). Additionally, in Equation 1, the symbol, " $\parallel$ " represents the concatenation of fields or functions.

$$H_{-}(L_0)=H(0\parallel H(L_0))$$

$$H_{-}(L_n)=H(H_{-}(L_{n-1})\parallel H(L_n))$$

EQUATION 1

As stated previously, the firmware measurement certificate **300** may be generated by a non-updateable, trusted piece of hardware or firmware code, such as the core root of trust, which is run at the first stage of initializing the component **214**. The core root of trust may measure the next layer or layers of firmware **216** by taking their cryptographic hash. Each measurement may contribute to the cumulative hash **308** and be included in the measurements **310**. In examples, the alias ID **312** may be generated by the core root of trust to authenticate the firmware measurement certificate **300**. More specifically, the alias ID **312** may be generated based on the cumulative hash **308** or the measurements **310**. The alias ID **312** may be used to sign the next layer. Accordingly, the alias ID **312** may be made available to the



next layer of firmware **216** within the responder once initialization is complete and used to digitally sign the firmware measurement certificate **300** for the subsequent next layer of the firmware **216**. For example, an alias ID, ID1, may be made available to layer L1. The layer, L1, may measure layer L2, which may generate alias ID, ID2, and issue a firmware measurement certificate **300** that is signed using the alias ID, ID1. This signature certifies the measurement of layer L2's cumulative hash **308** and measurements **310**, and alias ID2. The alias ID, ID2, may then be made available to layer L2.

FIG. 4 is an example system **400** for securing node groups. The system **400** includes multiple components, an initiator **402**, in communication with a responder **404**, over one or more interconnects **406**, and/or a network **408**. In examples, a controller, such as, the controller **212**, may initiate an authentication process for a component, such as the component **214**. Accordingly, the controller **212** may represent the initiator **402**, and the authenticated component **214** may represent the responder **404**. For example, the initiator **402** may be a general-purpose computer processor that performs in-band authentication of components **214** in a compute node **210** without a BMC. Alternatively, the initiator **402** may be a BMC, which performs out-of-band authentication of a general-purpose computer processor or other components **214** in a compute node with controllers **218**, such as a processor and a BMC. While the authentication performed in examples of the present disclosure may be performed against a variety of components **214**, for the sake of simplicity, authentication is described in the context of the authentication of a responder **404** that is a network interface controller (NIC). Accordingly, a BMC initiator and NIC responder may be connected to, and communicate over, a PCIe interconnect. Additionally, the NIC responder may be connected to a network **408**. The BMC initiator may be authenticating all the components **214** in the compute node **210** or performing an authentication on the NIC specifically in response to a call from the node group manager **204**. To determine whether the responder **404** is trustworthy, the initiator **402** may use an authentication service **428**. In examples, the authentication service **428** may be a computer application that uses digital certificates under a public key infrastructure (PKI) to determine the trustworthiness of the responder **404**. The responder **404** may not be trustworthy if a hacker or other malicious user has control. A malicious user may have control of the responder **404** if the responder **404** is a counterfeit hardware component, or if the firmware **216** on the responder **404** is counterfeit.

The initiator **402** and responder **404** may reside on the same compute nodes and thus communicate over one interconnect **406**, which passes messages between the initiator **402** and responder **404** based on a specific protocol. The interconnect **406** may include one or more interconnects, such as a USB, PCIe, Gen-Z, and the like. The initiator **402** and responder **404** may include protocol engines **426** that may ensure the messages between the initiator **402** and responder **404** are provided in a format that accords with the protocols of the relevant interconnects **406**. In some implementations, the initiator **402** may use multiple protocol engines **426** to handle interconnections with different kinds of components, such as, a baseboard management controller (BMC) and a general-purpose computer processor. The initiator **402** and responder **404** may also reside on different compute nodes **210**. In such a case, network components on the compute nodes may provide connection to the network **408**, which may include an internet protocol network, such as a local area network, wide area network, and the Internet.

To determine whether the responder **404** is trustworthy, the initiator **402** may authenticate the responder **404** by validating the public-private key pair **416** of the responder's component identification (ID) certificate **410** to determine the hardware of the responder **404** is authentic, i.e., not a counterfeit. The component ID certificate **410** may be a public key certificate that attests to the identity of the manufacturer of the responder **404**. Attesting to the identity of the responder's manufacturer may ensure that the hardware of the responder **404** is trustworthy. The component ID certificate **410** may be read by any entity wishing to authenticate the responder **404**, such as the initiator **402**. In examples, the initiator **402** may authenticate the responder **404** by identifying the public key of the public-private key pair **416** and determining if the responder **404** is in possession of the private key corresponding to the public key. If the responder **404** is in possession of the private key, the initiator **402** may determine that the responders' hardware is trustworthy. The public key may be identified by verifying that the component ID certificate **410**, or a chain of certificates, is signed by a trusted party. Once the public key may be trusted, the initiator **402** may challenge the responder **404** to prove possession of the corresponding private key. The challenge may involve having the responder **404** sign a nonce **438** with the private key. The nonce **438** may be a relatively large random number, e.g., 256 bits that is used only once. The initiator **402** may also apply an algorithm to the nonce **438** using the public key and use the resultant value to determine if the nonce **438** signed by the responder **404** is signed with the corresponding private key. If so, the initiator **402** may determine that the responder's hardware is trustworthy.

Additionally, the responder **404** includes firmware **216**, which may be a computer application that performs the operations of the responder **404**. For example, the NIC responder may operate a physical network, such as an Ethernet, wireless, or radio network. The NIC responder may also send and receive data packets from one compute node **210** to another. Another example responder **404** may be a disk controller. The disk controller responder **404** may read data from, and write data to, a hard disk drive according to a storage device protocol, such as serial advanced technology attachment (SATA). The firmware **216** may include one or more layers, wherein each layer represents one computer application that executes in a specified sequence. Thus, operation of the responder is performed by executing the layers of the firmware **216** in this sequence. In examples, determining whether the responder **404** is trustworthy may also involve determining whether the firmware **216** is trustworthy. In such examples, the initiator **402** may determine whether the firmware **216** is trustworthy by validating one or more firmware measurement certificates **414** for the firmware **216**.

As stated previously, the firmware measurement certificates **414** may be attribute certificates, which are digital documents that describe attributes that an issuer associates with the holder. In examples, the attributes described by the firmware measurement certificates **414** may be measurements of the binary image of the firmware **216** that is loaded into memory for execution. Attribute certificates may be associated with public key certificates, such as the component ID certificate **410**. In this way, the firmware measurement certificates **414** may extend the certificate chain used for authentication of the responder **404**. Thus, while the component ID certificate **410** may describe the identity of the manufacturer of the responder **404**, the firmware measurement certificates **414** may describe properties of the

firmware 216 being used to operate the responder 404. Similar to the component ID certificate 410, the firmware measurement certificates 414 may include a public key 432. The public key 432 may be part of a public-private key pair of mathematically-related keys for use in an asymmetric encryption scheme. In examples, the public-private keys pairs for each firmware measurement certificate 414 may be derived using an algorithm whose output depends on properties of the firmware or hardware logic that generates the certificate 414, such that a change in properties will produce a different pair of public-private keys. Additionally, the firmware measurement certificates 414 may include a cumulative hash 418, which may represent a measure of the firmware 216 when loaded into the computer memory (not shown) for execution. In examples, the cumulative hash 418 may be compared to binary image of expected firmware 422 stored on a known measurements server 434. In examples, the initiator 402 may cache measurements from the known measurements server 434 for comparison to the firmware measurement certificates 414. The expected firmware 422 may be a binary image of the firmware that is installed on the responder 404 during manufacture, or during a legitimate update from the manufacturer. Thus, if the cumulative hash 418 does not match a cumulative hash of the expected firmware 422, the firmware 216 may not be trustworthy. Accordingly, the initiator 402 may decline to use the responder 404. In examples, the firmware measurement certificates 414 may be issued by the component being verified, i.e., responder 404, during the process of initialization (power-up) and possibly at other times. Accordingly, the firmware measurement certificates 414 may reside in, and be retrieved from, the responder 404. Alternatively, the firmware measurement certificates 414 may be temporarily stored in memory for caching after being retrieved from the responder 404. Because the firmware measurement certificates 414 are signed, they can be cached safely. Any tampering would invalidate the signature and therefore the firmware measurement certificate 414.

Each layer of the firmware 216 may be associated with one of the firmware measurement certificates 414. In some examples, each firmware measurement certificate 414 may include a unique public key 432. In such examples, a chain of firmware measurement certificates 414 may be created between the different layers of the firmware 216. In other words, a layer n, of the firmware 216 may certify the public key of the next layer, layer n+1. In turn, the layer n+1 uses the private key associated with the certified public key to sign the firmware measurement certificate 414 of layer n+2. In other examples, a single private-public key pair may be used for all layers of the firmware 216 on the responder 404. In such examples, the different layers of firmware 216 may be chained together by updating a cumulative hash 418 of the firmware measurement certificate 414 for each layer of the firmware 216. Accordingly, to verify the link between two layers of the firmware 216, the initiator 402 may compare the cumulative hash 418 of each layer's firmware measurement certificate 414.

In order to ensure their trustworthiness, the firmware measurement certificates 414 may be generated by a core root of trust 424. The core root of trust 424 may include non-updateable hardware or firmware installed by the original manufacturer of the responder 404, which can be trusted to create the firmware measurement certificate 414 that represents the actual measured binary image of the firmware 216. In examples, the firmware 216 may include multiple layers. Each layer may represent one portion of the computer instructions used to operate the responder 404. The layers

may be executed in a prescribed sequence. Because the firmware 216 may include multiple layers, each layer may be susceptible to being compromised by a malicious user. Hence, the firmware measurement certificate 414 may be generated for each layer. To ensure the trustworthiness of the firmware measurement certificate 414 that is generated, the firmware measurement certificates 414 for each layer may be generated by an authenticated preceding layer. In an example responder 404 having multiple layers of firmware 216, the core root of trust 424 may generate a first firmware measurement certificate 414 representing the first layer of the firmware 216. Subsequently, prior to executing the second layer of firmware 216, the first layer may generate the firmware measurement certificate 414 for the second layer, thus ensuring the cumulative hash 418 for the second layer accurately represents the measured binary image of the second layer. Alternatively, the core root of trust 424 may generate a single firmware measurement certificate 414 that may be used to authenticate all the layers of the firmware 216.

Alternatively, a single firmware measurement certificate 414 may be used to authenticate multiple layers of firmware 414. Accordingly, a single firmware measurement certificate 414 may include measurements 436. The measurements 436 may represent a hash of the binary image for each layer of the firmware 216. Thus, there may be one measurement 436 for each layer of the firmware 216 up to the layer of the firmware 216 represented by the firmware measurement certificate 414. For example, if the firmware 216 includes layers, L0, L1, and L2, the measurements 436 for the firmware measurement certificate 414 may include three hashes: one hash for each of the binaries for layers, L0, L1, and L2.

Additionally, the firmware measurement certificate 414 may include a nonce 438. The nonce 438 may ensure the freshness of the measurements 436 and ensure that the core root of trust 424 is executed. The nonce 438 may be provided to the responder 404 by the initiator 402 during a challenge response protocol for authentication. Alternatively, the initiator 402 may write the nonce 438 to a specific memory location or register in the responder 404. Since the firmware measurement certificates 414 are generated at power-up or following a reset, the nonce 438 may be stored in a persistent location, such as in the responder 404. It is noted that for the first authentication of the responder 404, there may not be a nonce 438 available for the firmware measurement certificate 414. However, after the first authentication, the initiator 402 may provide the nonce 438, which may be written to persistent storage in the responder 404.

The responder 404 may include a signing service 430 and a protocol engine 426. The signing service 430 may provide a secure store. Examples of signing services include a trusted platform module and a field-programmable gate array. A trusted platform module may be a security co-processor that operates in response to a prescribed set of commands that may be used to securely store data, including the operating state of a computing platform, such as the compute node. A field-programmable gate array (FPGA) may be an integrated circuit that can be programmed using a hardware description language to perform specific instructions. In this way, an FPGA is similar to the processor. In contrast however, the processor may additionally be pre-programmed with a complex instruction set.

FIG. 5 is a process flow diagram of a method 500 for generating firmware measurement certificates to secure node groups. At block 502, a component of a compute node having firmware may power on or commence re-initializa-

## 11

tion. The component may be a responder, such as the responder **404**. An example responder **404** may include any of the components of the compute node **100**, such as the processor **108**, Gen-Z components **110**, and so on. For the purpose of this discussion, one of the PCIe components **130** of the compute node **100** is used as an example component. The PCIe component **130** may power on during a power-cycle of the compute node **100**. Alternatively, the PCIe component **130** may power on when the compute node **100** exits a low-power or reset state. A nonce, such as the nonce **438** may also be read from persistent storage. In examples, the nonce **438** may be provided by the initiator **402**. For example, the nonce **438** may be provided from a previous authentication request, e.g., authentication request **314**. In such an example, the responder **404** may store the nonce **438** in persistent storage within the responder **404**. In another example, the initiator **402** may write a nonce **438** into a register in the responder **404**. For example, PCIe devices, such as the PCIe component **130** may expose registers which can be read and written over a PCIe bus. The value stored in such a register may persist the nonce **438** across resets or power cycles

Additionally, the first layer, L0, of the firmware **412** of the PCIe component **130** may be loaded into memory. As stated previously, the PCIe component **130** may be an example of the responder **404** described with respect to FIG. 4. The PCIe component **130** may be powered on at block **502** by an operating system of the compute node **100**. The method **500** may be further performed by a core root of trust of the PCIe component **130**, such as the core root of trust **424**.

At block **504**, a hash may be generated for the next layer of the firmware **412** of the PCIe component **130**. For example, after powering on, the layer L0, which may represent the immutable core root of trust **424**, may measure layer L1 of the firmware **412**. The hash may be a NIST approved hash of the binary image of layer L1.

At block **506**, a firmware measurement certificate, such as the firmware measurement certificate **300**, may be generated for the next layer of the firmware **412**. For example, the core root of trust **424** may generate the firmware measurement certificate **300** for layer L1. The firmware measurement certificate **300** may include an issuer **302**, a component ID **306**, an alias ID **312**, and either a cumulative hash **308** or a measurement **310**. The issuer **302** and the component ID **306** may be the component ID of the PCIe component **130**. The component ID of the PCIe component **130** may be considered layer L0, which may be considered to be issuing the firmware measurement certificate **300**. The alias ID **312** may be a public key that identifies the owner of the firmware measurement certificate **300**. The owner of the firmware measurement certificate **300** may be the next layer, layer L1 after power on. The cumulative hash **308** or measurement **310** may be populated based on the hash generated at block **504**. The cumulative hash **308** may be determined based on Equation 1. Additionally, the measurement **310** may be populated with the generated hash. Additionally, the keyCertSign bit for the firmware measurement certificate **300** may be cleared to prevent malicious users from creating counterfeit certification authorities.

In examples, the firmware **412** may include one or more layers. As such, blocks **504-506** may be repeated for each subsequent layer of the firmware **412**. However, instead of being performed by the core root of trust **424**, the blocks **504-506** may be performed by the current layer of the firmware **412**. Thus, layer L0 may generate the firmware measurement certificate **300** for layer L1. Layer L1 may

## 12

generate the firmware measurement certificate for layer L2, and so on. If the firmware **412** includes one layer, the method **500** may flow to block **508**.

At block **508**, the layer of the firmware **412** associated with the generated firmware measurement certificate **300** may be executed. Executing the firmware **412** may involve operating the PCIe component **130**.

At block **510**, an initiator, such as the initiator **402**, may authenticate the firmware measurement certificate **300** generated at block **506**. The authentication may be performed as described with respect to FIG. 3. If the authentication fails, the method **500** may end. If the firmware measurement certificate **300** is authenticated however, the method **500** execution of the firmware **412** may continue. Further, in some scenarios, the same firmware measurement certificate **300** may be authenticated multiple times. Multiple authentications may be performed if the compute node **100** has not been power cycled, but authentication of the firmware **412** is requested. Additionally, the nonce provided by the initiator **402** during authentication may be written to the PCIe component **130**.

In examples, additional firmware measurement certificates **300** may be added after execution of the firmware **412**. If the operating system loads additional firmware **412** into the PCIe component **130** shortly after the OS has started, additional firmware measurement certificates **300** may be generated. Additional firmware measurement certificates **300** might also be generated if the operating system updates the firmware **412**. As such, the method may flow to block **504**. In the case of a single firmware measurement certificate **300** for all layers of the firmware **412**, the firmware measurement certificate **300** may be updated instead of adding new firmware measurement certificates **300**.

In some examples, multiple firmware measurement certificates **300** may be generated in blocks **504-510**, one for each layer of the firmware **412**. In such examples, the component ID **306** for each firmware measurement certificate **300** may be the component ID of the PCIe component **130**. The issuer **302** for such firmware measurement certificates **300** may be the preceding layer of the firmware **412**. The alias ID **312** may be a public key that identifies the current layer of the firmware **412**. The nonce **438** may be the nonce stored in the PCIe component **130** by the initiator **402**. In examples with multiple firmware measurement certificates **300** for one component, such as the PCIe component **130**, the cumulative hash **308** may be populated. The cumulative hash **308**, as described with respect to Equation 1, may be a concatenation of the previous layer's cumulative hash and an NIST-approved hash function of the binary image of the current layer. The NIST-approved hash function of the binary image of the current layer may additionally be appended to the measurements **310**.

In some examples, a single firmware measurement certificate **300** may be generated that represents all the layers of the firmware **412**. In such examples, instead of generating new firmware measurement certificates **300** for each layer, a new firmware measurement certificate **300** may be issued with an updated signature. This new firmware measurement certificate **300** may replace the previously issued firmware measurement certificate **300**. This new firmware measurement certificate **300** may be generated by the current layer of the firmware **412**. In the new firmware measurement certificate **300**, the component ID **306** and nonce **314** may not be changed. However, the issuer **302** may be the current layer of firmware **412**. The alias ID **312** may be a public key that identifies the next layer of the firmware **412**. In examples with a single firmware measurement certificate

13

300 for the PCIe component 130, the cumulative hash 308 and the measurements 310 may be populated. The cumulative hash 308, as described with respect to Equation 1, may be a concatenation of the previous layer's cumulative hash and an NIST-approved hash function of the binary image of the current layer. Accordingly, the measurements 310 from the previously issued firmware measurement certificate 300 may be supplemented with a new measurement 310 for the next layer of the firmware 412. The measurement 310 may be an NIST-approved hash function of the binary image of the next layer of firmware 412.

It is to be understood that the process flow diagram of FIG. 5 is not intended to indicate that the method 500 is to include all of the blocks shown in FIG. 5 in every case. Further, any number of additional blocks can be included within the method 500, depending on the details of the specific implementation. In addition, it is to be understood that the process flow diagram of FIG. 5 is not intended to indicate that the method 500 is only to proceed in the order indicated by the blocks shown in FIG. 5 in every case. For example, block 504 can be rearranged to occur before block 502.

FIG. 6 is a message flow diagram 600 for authenticating components in a system for securing node groups. The message flow diagram 600 may represent a flow of messages between an authentication initiator 602 and a responder 604. The initiator 602 may represent a component, such as the initiator 402, and may include an authentication service 606 and a protocol engine 608. The protocol engine 608 may translate the messages between the initiator 602 and the responder 604 based on the interconnects between the initiator 602 and responder 604. The responder 604 may represent a component having firmware, such as the responder 404. Message 610 represents a request from the initiator 602 for the responder's certificate chain (or chains.) The message 612 represents the certificate chain sent by the responder 604 to the initiator 602 in response to the request. In response to receiving the responder's certificate chain, the initiator 602 may validate one or more of the certificate chains and select the public key to be authenticated by the responder 604. The public key may be selected from the leaf certificate of a valid certificate chain.

Message 614 may represent an authentication request from the initiator 602 to the responder 604. The authentication request may consist of a large random nonce, and the selected public key to be authenticated. In the general case, the public key to be authenticated is identified because the responder 604 may have multiple public-private key pairs used for different purposes.

Once the authentication request is received, the responder's protocol engine may extract the nonce and identity of the public key to be authenticated from the authentication request. Additionally, the responder's signing service may sign a concatenation of the nonce and an internally generated random salt with the private key that corresponds to the identified public key. The purpose of the salt is to protect against chosen plain-text attacks, so it should be unpredictable to the initiator. Message 616 may represent the responder's response to the authentication request containing the salt and the signature over the concatenation of the nonce and salt.

Once the response to the authentication request is received, the initiator 602 may verify that the nonce and salt have been signed by the private key corresponding to the public key in the leaf certificate. If verification is successful, the responder 604 has been authenticated.

14

FIG. 7 is an example system 700 for securing node groups. The example system 700 includes a node group manager 702, a baseboard management controller (BMC) 704, a processor 706, and a known measurements server 708. The BMC 704 and processor 706 may represent specific examples of the controller 212 described with respect to FIG. 2. Additionally, the BMC 704 and processor 706 may be implemented on one or more nodes 210 within a node group (not shown). Similar to the system 200, the example node group manager 702 may secure one or more node groups by authenticating each of the nodes 210. Authenticating each of the nodes 210 may include authenticating the trustworthiness of the hardware and firmware of the BMC 704, processor 706, and components 712 that are installed. In examples, the components 712 may be connected to either or both of the BMC 704 and processor 706 over one or more interconnects, including, PCIe, USB, I2C, and the like. Additionally, the node group manager 702 may be used to report the inventory status of each of the nodes over an IP network. The inventory status may represent various details about the authentication of an inventory of components 712 for the node 210. The inventory status may include an identifier of the BMC 704, processor 706, or component 712 being authenticated, the timing of the authentication(s), the interconnect(s) over which the authentication is performed, and the like. Being able to determine the inventory status of the components 712 may make it possible to detect and monitor activity of service engineers, for example. In this way, it may be possible to reduce the likelihood of mistakes or malicious insertion during service activity. Service activity refers to the acts of installing new components 712 or updating firmware.

The BMC 704 and processor 706 may include protocol engines 714, an authentication service 716, and a component authentication manager 718. The authentication service 714 and protocol engines 716 may represent examples of the authentication service 606 and protocol engine 608 described with respect to FIG. 6. Referring back to FIG. 7, the component authentication manager 718 may use the authentication service 714 and protocol engines 716 to perform the authentications during initialization and at the request of the node group manager 702.

During initialization, the component authentication manager 718 may direct the authentication service 714 to authenticate every component 712 that is reachable on each interconnect or fabric that connects to the respective BMC 704 and processor 706. Some components 712 may be accessible by more than one interconnect. For example, a component 712 may be accessible to the BMC 704 over both PCIe and I2C. In such an example, a point-to-point I2C connection may be used to distribute confidential material, e.g. encryption or authentication keys. The I2C connection may be used in this way instead of a shared PCIe fabric. The shared PCIe fabric may have relatively higher bandwidth than the I2C interconnect and may be used for data transfers. However, the PCIe fabric may be more susceptible to snooping than I2C. Thus, by authenticating the component 712 through both its PCIe and its I2C connections, the BMC 704 may have assurance that the component 712 is securely connected. Keys distributed over the I2C connection may be used to encrypt and decrypt data that is stored within a compute node. Alternatively, encryption keys may be used for authenticated and encrypted data transfers for data that is stored in the compute node unencrypted.

By authenticating the components 712 over the available interconnects and fabrics, the component authentication manager 718 may generate an authenticated inventory of all

15

active components **712** to which the respective BMC **704** or processor **706** is connected. More specifically, the component authentication manager **718** may include a data store having a list of authenticated components, together with their certificate chains and the status of whether each certificate chain has been authenticated or not. Generating and maintaining this authenticated inventory over time may provide a record of inventory tracking that can be validated against the known measurements server **708** to ensure the trustworthiness of the components **712** and associated firmware. The component authentication manager **718** may also be able to build and maintain a connection topology for the components **712**. Additionally, the BMC **704** may be connected over an interconnect or fabric with the processor **706**. As such, the BMC **704** and processor **706** may authenticate each other.

In examples, the BMC **704** may serve as the local authority for all communication with the node group manager **702**. Although FIG. 7 shows a single BMC **704** connected to the node group manager **702**, in examples, the node group manager **702** may maintain connections to several BMCs. Similarly, although the system **700** only includes one node group manager **702**, in examples, a hierarchy, federation, or other set of cluster authentication and verification managers may be present. The component authentication manager **718** in the BMC **704** may retrieve the authenticated inventory from the component authentication manager **718** in the processor **706**. In examples, the authenticated inventory may be stored in an attribute certificate that contains a list of component identities that have been authenticated by the controller, along with a set of metadata, e.g., interconnect identifier, timestamp, etc. Such an attribute certificate may also list the authentication failures (e.g., identifiers for the components **712** that failed to authenticate, were unresponsive, or had unresponsive connections). The attribute certificate may be issued by the processor **706** once the processor **706** has been authenticated to the BMC **704**. Since attribute certificates may be signed assertions, the attribute certificates may be stored, copied and transmitted over networks without danger that a malicious user may add or remove information about what components **712** have been authenticated.

Authentication failures may be detected through electrical connections with unresponsive components **712**, missing authentications from components **712** that were previously authenticated, and changes to the connection topology for components **712**. Accordingly, the component authentication manager **718** and node group manager **702** may store previous connection topologies, inventory reports, and authentication status for comparison.

The node group manager **702** may authenticate the BMC **704** by using the BMC's component authentication manager **718** to retrieve an attribute certificate including the authenticated inventory of components **712** authenticated by the BMC **704**. In addition, the node group manager **702** may use the BMC's component authentication manager **718** to retrieve the authenticated inventory of the processor **706**. In an example, the inventory of components for a node, such as the node **210**, may be included in a single signed assertion or certificate. In such examples, the certificate or assertion may reference other certificates and assertions.

The node group manager **702** may also be able to retrieve all certificate chains for all components **712** using the component authentication managers **718** and may be able to request that the component authentication managers **718** authenticate any certificate chain for any connected component **712** or controller. In the case of components connected to the

16

processor **706**, the BMC's component authentication manager **718** may forward the authentication request to the processor's component authentication manager **718**.

In an example, the component authentication manager **718** may provide an application programming interface (API). In such an example, the API may include operations to list components, get certificate chains, query status, authenticate components and other related operations. The list components operation may return one or more attribute certificates. Each attribute certificate may list the authenticated components **712** associated with a particular component authentication manager **718**, as well as relevant metadata such as interconnect identifiers and any authentication failures that may have been detected. In examples, component authentication managers **718** may be logically configured in a hierarchy. Accordingly, certain component authentication managers **718** may represent different levels of an architecture in a tree-like structure. As such, the component authentication managers **718** at the bottom of this tree-like structure may be referred to as leaves. In examples, the component authentication managers **718** for leaves in the hierarchy may return a single attribute certificate. For component managers **718** higher in the hierarchy, multiple attribute certificates will be returned. In examples, variations of the list component operation may also be included. Such variations may provide information about the topology of the component **712**, and authentication information, such as a timestamp of the authentication.

The get certificate chains operation may provide the certificate chains retrieved from a component **712**. Accordingly, a call to the get certificate chains operation may specify the component **712** for which the certificate chains are to be retrieved.

The query status operation may provide the authentication status of the specified component **712**. If the component **712** has been authenticated, the query status operation may provide the public-private key pair **416** that has been authenticated. Additionally, for each public-private key pair **416**, the query status operation may provide the number of times that the public-private key pair **416** has been authenticated since the component **712** was last initialized and a timestamp of the last authentication. Further, for authenticated components **712**, the identity of the authentication initiator may be provided. For components **712** that have failed authentication, an error code may be provided.

The authenticate operation may specify an identifier for the component **712** being authenticated, a public key of the public-private key pair **416** for the component **712**, and a nonce. Accordingly, the component authentication manager **718** may use the given nonce to authenticate the identified component **712** that is in possession of the given public key. However, in some cases, the authenticate operation may be denied if it is not possible to authenticate the component **712**. For example, it might not be possible for the component authentication manager **718** of the processor **706** to authenticate a DDR component if the processor **706** is running a workload, such as applications or the operating system. Rather, such authentication may be limited to times when the processor **706** is in a particular mode of operation, or running specific firmware: for example, during a power-on-self-test.

Using firmware measurement certificates **218** to denote authenticated inventory may prevent tampering because the inventory can be compared over time. For example, each time a node is rebooted, the component authentication manager **718** may compare the authenticated inventory to the actual inventory. Once the node group manager **702**

accesses the firmware measurement certificates **218** of authenticated components **712** and the certificate chains corresponding to the authenticated components **712**, the component authentication manager **718** may verify the correctness of any firmware measurements contained in authenticated certificate chains. The component authentication manager **718** may verify the correctness by comparing the measurements or cumulative hashes to values contained in the known measurements server **708**.

How the authentication of components **712** is done may vary depending on whether the system **700** is undergoing initialization or running a workload. After a reset or power-cycle, as part of the initialization process, the component authentication manager **718** may authenticate all components **712** on any bus or fabric connected to the BMC **704** described above. The results of the authentication may be saved in the form of an attribute certificate in the authentication database **206**. Additionally, the BMC **704** may boot the processor **706** with firmware or a reduced-function operating system so that the processor **706** can authenticate all components **712** on any bus or fabric connected to the processor **706**, such as the DDR, Gen-Z, PCIe and USB connections. Additionally, the BMC **704** may authenticate the processor **706** using the authentication pattern described with respect to FIG. 6. In examples, the system **700** may include a security co-processor (not shown), such as a trusted platform module (TPM). In such examples, the processor **706** may authenticate the keys and certificates that have been issued to the TPM during manufacturing. The TPM may be attached to the same board as the processor **706**. Alternatively, the TPM may be integrated within the processor package. In examples, the BMC **704** may retrieve an attribute certificate (not shown) from the processor's component authentication manager **718** that lists the authenticated inventory of components **712** connected to the processor **706**. This attribute certificate may be made available to the node group manager **702** where it may be archived for future comparison and inventory tracking. In this way, all components **712** within the system **700** may be authenticated during power-cycle and initialization.

However, the system **700** may not be power-cycled or re-initialized for months or years. Hence, during runtime, the BMC **704** may run the component authentication manager **718** to authenticate all connected components **712**. Additionally, the BMC **704** may authenticate any new component **712** that is the subject of a hot-plug or hot-insertion event, such as a USB device. In examples, the operating system and device drivers may be enhanced to support component authentication of the processor **706** and the connected components. In such examples, the operating system may run an authentication service and also the component authentication manager **718** to report the authentication status of any components **712** that are authenticated. The operating system may maintain presence, insertion and removal information for all connected components **712**. Further, the operating system may authenticate components **712** in response to insertion events and periodically during runtime. The operating system may also authenticate the processor **706** and connected components **712** when an exit from a low power state occurs.

Alternatively, a security processor distinct from the processor **706** and a Trusted Execution Environment (TEE) may be used to run the component authentication manager **718** and authentication service. In another example, a reduced-function hypervisor may be used to wall-off component authentication and other security critical functionality from the main operating system. In another example, a system

management interrupt (SMI) may be used to run the component authentication manager **718** and authentication service on the processor **706**. In such an example, the BMC **704** may generate the SMI to trigger the execution of a handler that runs the component authentication manager **718** and authentication service. In this way, the BMC **704** may trigger an authentication of one or more components **712** connected to the processor **706**. In some scenarios, the components **712** may be subject to deep low-power events. In such scenarios, runtime authentication may be used to detect any component replacement that occurs during such events.

FIG. 8 is an example system **800** for securing node groups. Other arrangements of the authentication are possible and may be more useful for different architectures. For example, system **800** includes a cluster authentication and verification manager **802**, a rack authentication and verification manager **804**, a processor **806**, and known good measurements **808**. The cluster authentication and verification manager **802** and rack authentication and verification manager **804** may be similar to the node group manager **702** described with respect to FIG. 7 but operating at different hierarchical levels to address manageability and scalability requirements. Referring back to FIG. 8, the processor **806** may include protocol engines **814**, an authentication service **816**, and a component authentication manager **818**. The authentication service **814** and protocol engines **816** may represent examples of the authentication service **606** and protocol engine **608** described with respect to FIG. 6. Referring back to FIG. 8, the component authentication manager **818** may use the authentication service **814** and protocol engines **816** to perform the authentications during initialization and at the request of the cluster authentication and verification manager **802**.

The system **800** may represent a node without a BMC where a single component authentication manager **818** runs on the processor **806** and is connected to the rack authentication and verification manager **804**. The system **800** may perform "in-band" authentication, which means the authentication may be taking place over the same interconnections used by the applications and services that are running on the processor **806**. In examples, the rack authentication and verification manager **804** may be connected to multiple processors, e.g., all the processors in a rack. The rack authentication and verification manager **804** may perform the same functions as the cluster manager, such as, verifying the firmware, but only for those components **812** in the rack. In addition, the rack authentication and verification manager **804** may provide the same API to the cluster authentication and verification manager **802** that the component authentication manager **718** provides to the node group manager **702** as described with respect to FIG. 7. In this way, the cluster authentication and verification manager **802** may be able to get a cluster-wide view of the authenticated inventory and force re-authentication of any component **812** through the rack authentication and verification manager **804**. In another example, the system **800** may include a dedicated network interface for the management plane, wherein access to the component authentication manager **818** may be restricted to the dedicated network interface.

In another example, some components **812** may not be associated with a node, rather such components **812** may be associated with a rack or another component, such as a power distribution unit, fan, or sensor. In such an example, the components **812** may be authenticated using the same mechanisms and principles described above. Accordingly, the component authentication manager **818** and authentication service **814** may be embedded in the rack authentication

and verification manager **804** and use whatever protocol is used to control the rack power distribution unit, fan or sensor for authentication.

In examples, the known good measurements **808** may include manifests that described the expected measurements of the components **812**. The manifests may be created by manufacturers of platforms, such as nodes **210**. More specifically, during the last stages of manufacture of a node **210**, e.g., during testing, the components **812** could be powered-up and initialized. Further, firmware measurement certificates may be generated for the components. Accordingly, the cluster authentication and verification manager **802** and the rack authentication and verification manager **804** may determine the component identities and firmware measurements of the components **812**. The manufacturer or system integrator can then issue a set of “manifest” certificates specifying what components and firmware are in the node **210**, i.e. what is the “known good state” of the overall system as it exits the factory. Thus, when the node **210** arrives at the customer site, the customer can compare the manifest to what is being reported by the cluster authentication and verification manager **802** or the rack authentication and verification manager **804** to ensure that the node **210** has not been tampered with en route from the factory. In one example, the manifests may be Platform Certificates, which are under development by the Trusted Computing Group.

With regard to authentication of node groups, a hierarchy of services may be used to manage component authentication and verification at cluster scale. Accordingly, Table 1 provides an example hierarchy listing the services and responsibilities of each node hierarchy level in the hierarchy. The lowest level in the hierarchy is a single node, whereas the highest level may be the cluster. The node may deliver compute, storage or networking (switch) services. The bulk of components may be contained within nodes. An enclosure may contain multiple nodes and may have one or more power supplies. A rack may have multiple enclosures and power supplies or UPS (Uninterruptable Power Supplies). A cluster may contain multiple racks.

TABLE 1

Node Hierarchy Level	Services and responsibilities
Single Node	Component authentication manager and component authentication service Services may run on iLO, BMC, Management Engine, and main processor Authenticate all components on the node Store authenticated firmware measurement certificates Periodic re-authentication of components Provide authenticated inventory in the form of attribute certificates to higher level entities Respond to authentication requests from higher level entities Authenticate self, authenticate specified component Enforce authentication and verification policy
Enclosure	Component authentication manager, component authentication service, verification service Services may run on designated management node Authenticate all components in the enclosure Directly for components not associated with a node Via component authentication manager for components within a node Request authenticated inventory from node component authentication managers Compares firmware measurements to known good state Recovery action initiated if measurements determine to be incorrect

TABLE 1-continued

Node Hierarchy Level	Services and responsibilities
Rack	Store authenticated certificates Periodic re-authentication of components via component authentication managers Provide authenticated inventory in the form of attribute certificates to higher level entities Respond to authentication requests from higher level entities Authenticate self, authenticate specified component via component authentication managers Enforce authentication and verification policy Component authentication manager, component authentication service, verification service Services may run on designated management node Authenticate all components in the rack Directly for components not associated with an enclosure Via component authentication manager for components within an enclosure Request authenticated inventory from enclosure component authentication managers Compares firmware measurements to known good state Recovery action initiated if measurements determine to be incorrect Store authenticated certificates Periodic re-authentication of components Provide authenticated inventory in the form of attribute certificates to higher level entities Respond to authentication requests from higher level entities Authenticate self, authenticate specified components via component authentication managers Enforce authentication and verification policy
Cluster	This level of hierarchy may use multiple rack, enclosure and node, authentication and verification services to manage authentication and verification at cluster scale. Various architectures for high availability and resilience may be employed.

Depending on the size and configuration of a computer hardware architecture, some elements of the hierarchy may be omitted. For example, if there are a small number of racks, the cluster services may interact directly with the enclosure authentication and verification services. It is noted that Table 1 is merely one possible hierarchy of nodes and node groups. In examples, other organizations of hierarchies may be used. For example, clusters may be included within the hierarchy level of a data center. Further, data centers may be included within a hierarchy level of an information technology department, or an enterprise, for example.

FIG. 9 is a process flow diagram of a method **900** for securing node groups. The method **900** may be performed by a node group manager, such as the node group manager **204**.

At block **902**, the node group manager **904** may authenticate a hardware architecture of the components **214** of a node group **202**. The node group **202** may be described as a node hierarchy level. The node hierarchy level may include one or more compute nodes **210**. Authenticating the hardware architecture may include establishing that the hardware of the components **214** is from known and trusted manufacturers. Establishing that the hardware is from known and trusted manufacturers may include validating the certificate chain of the components **214**.

At block **904**, the node group manager **204** may authenticate the firmware **216** of each of the components **214** of the node group **202**. Authenticating the firmware **216** may include comparing the measurements in the firmware measurement certificates **218** for the components **214** to the measurements in the known measurements server **208**.

At block **906**, the node group manager **204** may generate an authentication database **206** based on the authenticated

21

hardware architecture and the authenticated firmware **216**. The authentication database **206** may include descriptions of the authentications, such as when each component **214** is authenticated.

At block **908**, the node group manager **204** may secure the node group **210** by using the authentication database **206**. In examples, a policy for securing a specified node group **202** may be implemented by using the authentication database **206**. For example, the authentication database **206** may describe when the components **214** of the node group **202** were last authenticated. Further, an example policy for securing the node group **202** may specify that the components **214** are to be authenticated at least once a month. As such, the node group manager **204** may execute a policy script every month that checks when the components **214** of the node group **202** were last authenticated. If any of the components **214** has not been authenticated within the last month, the node group manager **204** may automatically validate the components **214** that have not been validate in accord with the policy.

It is to be understood that the process flow diagram of FIG. **9** is not intended to indicate that the method **900** is to include all of the blocks shown in FIG. **9** in every case. Further, any number of additional blocks can be included within the method **900**, depending on the details of the specific implementation. In addition, it is to be understood that the process flow diagram of FIG. **9** is not intended to indicate that the method **900** is only to proceed in the order indicated by the blocks shown in FIG. **9** in every case. For example, block **904** can be rearranged to occur before block **902**.

FIG. **10** is an example system **1000** comprising a tangible, non-transitory computer-readable medium **1002** that stores code for securing node groups. The tangible, non-transitory computer-readable medium is generally referred to by the reference number **1002**. The tangible, non-transitory computer-readable medium **1002** may correspond to any typical computer memory that stores computer-implemented instructions, such as programming code or the like. For example, the tangible, non-transitory computer-readable medium **1002** may include RAM, ROM, EEPROM, CD-ROM or other optical disk storage, magnetic disk storage or other magnetic storage components, or any other medium that may be used to carry or store desired program code in the form of instructions or data structures and that may be accessed by a computer. Disk and disc, as used herein, includes compact disc (CD), laser disc, optical disc, digital versatile disc (DVD), floppy disk and Blu-ray® disc where disks usually reproduce data magnetically, while discs reproduce data optically with lasers.

The tangible, non-transitory computer-readable medium **1002** can be accessed by a processor **1004** over a computer bus **1006**. The processor **1004** may be a central processing unit that is to execute an operating system in the system **1000**. A region **1008** of the tangible, non-transitory computer-readable medium **1002** stores computer-executable instructions that authenticate a hardware architecture of each of a plurality of components of the compute nodes. A region **1010** of the tangible, non-transitory computer-readable medium stores computer-executable instructions that authenticate a firmware of each of the components. A region **1012** of the tangible, non-transitory computer-readable medium stores computer-executable instructions that generate an authentication database comprising a plurality of authentication descriptions that are based on the authenticated hardware architecture and the authenticated firmware, wherein a policy for securing a specified subset of the

22

compute nodes is implemented by using the authentication database. A region **1014** of the tangible, non-transitory computer-readable medium stores computer-executable instructions that secure the node group by using the authentication database.

Although shown as contiguous blocks, the software components can be stored in any order or configuration. For example, if the tangible, non-transitory computer-readable medium **1002** is a hard drive, the software components can be stored in non-contiguous, or even overlapping, sectors.

The foregoing description, for purposes of explanation, used specific nomenclature to provide a thorough understanding of the disclosure. However, it will be apparent to one skilled in the art that the specific details are not required in order to practice the systems and methods described herein. The foregoing descriptions of specific examples are presented for purposes of illustration and description. They are not intended to be exhaustive of or to limit this disclosure to the precise forms described. Obviously, many modifications and variations are possible in view of the above teachings. The examples are shown and described in order to best explain the principles of this disclosure and practical applications, to thereby enable others skilled in the art to best utilize this disclosure and various examples with various modifications as are suited to the particular use contemplated. It is intended that the scope of this disclosure be defined by the claims and their equivalents below.

What is claimed is:

1. A method comprising:

accessing, by a node group manager, data representing an inventory of components of a plurality of computer nodes and accessing details about prior authentications of the inventory of components;

verifying, by the node group manager, compliances of the components based on the details and a security policy, wherein verifying the compliances comprises determining a given component of the inventory of components does not comply with the security policy; and responsive to determining that the given component does not comply with the security policy, performing, by the node group manager, a responsive action.

2. The method of claim 1, wherein:

the given component is associated with a given prior authentication of the prior authentications;

a given detail of the details represents a time associated with the given prior authentication; and determining that the given prior authentication violates the security policy comprises determining that the time is not current according to the security policy.

3. The method of claim 1, wherein:

the given component is associated with a given prior authentication of the prior authentications;

a given detail of the details represents that an interconnect associated with the given prior authentication is different from an interconnect currently being used by the given component; and

performing the responsive action comprises performing an action to quarantine the given component.

4. The method of claim 1, wherein:

the given component is associated with a given prior authentication of the prior authentications;

the given prior authentication comprises executing firmware over first interconnect;

a given detail of the details represents that the first interconnect is different from a second interconnect currently being used by the given component to execute firmware; and



23

performing the responsive action comprises performing an action to quarantine the given component.

5. The method of claim 1, wherein the given component is part of a given node of the plurality of nodes, wherein the given node comprises a controller, and wherein performing the responsive action comprises authenticating, by the controller, the given component.

6. The method of claim 5, wherein the authenticating comprises providing, by the node group manager, a call to an application programming interface (API) of the controller.

7. The method of claim 1, wherein the given component is part of a given node of the plurality of nodes, wherein the given node comprises a controller, and wherein accessing the data comprises providing, by the controller, an attribute certificate comprising data representing a list of components of the inventory of components which are part of the given node.

8. The method of claim 7, wherein the attribute certificate further comprises data representing an additional component of the inventory of components which failed authentication.

9. The method of claim 1, wherein accessing the data comprises accessing, by the node group manager, an authentication database.

10. The method of claim 1, wherein determining that the given component does not comply with the security policy comprises determining, by the node group manager, that authentication of firmware of the given component does not comply with the security policy.

11. The method of claim 1, wherein determining that the given component does not comply with the security policy comprises determining, by the node group manager, that a hardware architecture associated with the given component does not comply with the security policy.

12. A system comprising:  
a node group comprising a plurality of computer nodes, wherein the plurality of computer nodes comprises a plurality of components;  
a data store associated with data representing details about prior authentications of the inventory of components; and  
a node group manager comprising a hardware processor to:  
access the data associated with the data store;  
responsive to the access, verify compliances of components of the inventory of components based on the details and a security policy, wherein verifying the compliances comprises determining that a given component of the inventory of components does not comply with the security policy; and  
perform a responsive action responsive to the determination.

13. The system of claim 12, wherein:  
the given component is associated with a given prior authentication of the prior authentications; and  
the hardware processor to further:  
determine that a given detail of the details represents that an interconnect associated with the given prior authentication is different from an interconnect currently being used by the given component; and  
perform the responsive action to quarantine the given component.

24

14. The system of claim 12, wherein:

the given component is associated with a given prior authentication of the prior authentications;

a given detail of the details represents a time associated with the given prior authentication; and

the hardware processor to further determine that the given prior authentication violates the security policy based on the time not being current according to the security policy.

15. The system of claim 12, wherein the data store comprises an authentication database.

16. A non-transitory storage medium to store machine-readable instructions that, when executed by a machine associated with a node group manager, cause the machine to:  
access data representing an inventory of components of a plurality of computer nodes and accessing details about prior authentications of the inventory of components;  
verify compliances of the components based on the details and a security policy, wherein verifying the compliances comprises determining a given component of the inventory of components does not comply with the security policy; and

responsive to determining that the given component does not comply with the security policy, perform a responsive action.

17. The storage medium of claim 16, wherein:  
the given component is associated with a given prior authentication of the prior authentications;  
a given detail of the details represents a time associated with the given prior authentication; and  
the instructions, when executed by the machine, further cause the machine to:  
determine that the time is not current according to the security policy; and  
responsive to the determination that the time is not current, authenticate the given component.

18. The storage medium of claim 16, wherein:  
the given component is associated with a given prior authentication of the prior authentications;  
the given prior authentication comprises executing firmware over a first interconnect; and  
the instructions, when executed by the machine, further cause the machine to:  
determine, based on a given detail of the details, that the first interconnect is different from a second interconnect currently being used by the given component to execute firmware; and  
responsive to determining that the first interconnect is different from the second interconnect, perform an action to quarantine the given component.

19. The storage medium of claim 16, wherein the instructions, when executed by the machine, further cause the machine to determine, based on the data, that authentication of firmware of the given component does not comply with the security policy.

20. The storage medium of claim 16, wherein the instructions, when executed by the machine, further cause the machine to determine, based on the data, that authentication of a hardware architecture of the given component does not comply with the security policy.

\* \* \* \* \*