

# US Patent & Trademark Office

## Patent Public Search | Text View

---

|                      |                              |
|----------------------|------------------------------|
| United States Patent | 12393724                     |
| Kind Code            | B2                           |
| Date of Patent       | August 19, 2025              |
| Inventor(s)          | Remseth; Bjorn Gunnar et al. |

---

### Privacy aware source code

---

#### Abstract

Systems and methods for protecting privacy-relevant data from unauthorized disclosure in source code of an application. For instance, the present disclosure provides a plurality of technical features including: a privacy-relevant data analyzer that analyzes source code, detects privacy-relevant data in the source code, and generates a report of instances of detected privacy-relevant data. In some examples, the privacy-relevant data analyzer scans through source code to detect annotations that denote if fields, records, or combinations thereof include privacy-relevant data. The privacy-relevant data analyzer further generates and provides a report of detected privacy issues associated with sensitive data included in source code so that the issues can be resolved to ensure that privacy is not breached.

---

**Inventors:** Remseth; Bjorn Gunnar (Viken, NO), Chen; Haiyue (Oslo, NO), Reißmann; Nico (Trondheim, NO), Skogvold; Arve Aleksander Nymo (Trondheim, NO)

**Applicant:** Microsoft Technology Licensing, LLC (Redmond, WA)

**Family ID:** 1000008766230

**Assignee:** Microsoft Technology Licensing, LLC (N/A, N/A)

**Appl. No.:** 18/048953

**Filed:** October 24, 2022

#### Prior Publication Data

| Document Identifier | Publication Date |
|---------------------|------------------|
| US 20240135030 A1   | Apr. 25, 2024    |
| US 20240232424 A9   | Jul. 11, 2024    |

---

#### Publication Classification

**Int. Cl.:** G06F21/62 (20130101)

**U.S. Cl.:**

CPC                    **G06F21/6245** (20130101);

## Field of Classification Search

**CPC:**                G06F (21/6245)

---

## References Cited

### U.S. PATENT DOCUMENTS

| Patent No.   | Issued Date | Patentee Name   | U.S. Cl. | CPC          |
|--------------|-------------|-----------------|----------|--------------|
| 10180836     | 12/2018     | Arguelles       | N/A      | G06F 8/75    |
| 10635434     | 12/2019     | Greer           | N/A      | N/A          |
| 11182490     | 12/2020     | Allen et al.    | N/A      | N/A          |
| 11941156     | 12/2023     | Kats            | N/A      | H04L 63/102  |
| 11973784     | 12/2023     | Erlingsson      | N/A      | G06F 9/5077  |
| 12058160     | 12/2023     | Erlingsson      | N/A      | G06F 16/906  |
| 12093414     | 12/2023     | Munoz           | N/A      | G06F 21/6218 |
| 12126643     | 12/2023     | Skarphedinsson  | N/A      | G06F 21/552  |
| 2019/0213355 | 12/2018     | Raviv           | N/A      | G06F 16/901  |
| 2021/0011919 | 12/2020     | Lu et al.       | N/A      | N/A          |
| 2022/0229856 | 12/2021     | Malhotra        | N/A      | G06F 21/577  |
| 2023/0161585 | 12/2022     | Medeiros Destro | 717/123  | G06F 8/73    |
| 2023/0305813 | 12/2022     | Jalal           | N/A      | G06F 9/455   |
| 2024/0078096 | 12/2023     | Gass            | N/A      | G06F 8/65    |
| 2024/0078336 | 12/2023     | Sansom          | N/A      | G06F 8/77    |
| 2024/0086570 | 12/2023     | Sharma          | N/A      | G06F 21/6245 |
| 2024/0095027 | 12/2023     | Fanning         | N/A      | G06F 8/77    |
| 2024/0135030 | 12/2023     | Remseth         | N/A      | G06F 21/6245 |
| 2024/0135477 | 12/2023     | Fron            | N/A      | G06F 21/54   |
| 2024/0211232 | 12/2023     | Gass            | N/A      | G06F 8/40    |
| 2024/0242158 | 12/2023     | Hora            | N/A      | G06Q 10/105  |
| 2024/0256254 | 12/2023     | Gass            | N/A      | G06F 8/65    |
| 2024/0311476 | 12/2023     | Cohen           | N/A      | G06F 21/53   |
| 2024/0345993 | 12/2023     | Madisetti       | N/A      | G06F 16/176  |
| 2024/0354408 | 12/2023     | Mooney, III     | N/A      | G06F 21/563  |

### FOREIGN PATENT DOCUMENTS

| Patent No. | Application Date | Country | CPC |
|------------|------------------|---------|-----|
| 111709052  | 12/2020          | CN      | N/A |

### OTHER PUBLICATIONS

Jaager, et al., “PrivaLog: A Privacy-Aware Logic Programming Language”, In Repository of arXiv:2010.07057v2, May 17, 2021, 41 Pages. cited by applicant

Jelínek, Vojtěch, “Privacy Aware Programming”, Retrieved from:

<https://github.com/oppia/oppia/wiki/Privacy-aware-programming>, Oct. 15, 2021, 8 Pages. cited by applicant

Li, et al., “Honeysuckle: Annotation-Guided Code Generation of In-App Privacy Notices”, In

*Primary Examiner:* Gundry; Stephen T

---

## **Background/Summary**

### **BACKGROUND**

(1) Source code for an application can include and/or process sensitive data, such as Personally Identifiable Information (PII), attorney-client privileged information, Protected Health Information (PHI), or other information that an individual or enterprise wants to keep from unauthorized disclosure. For instance, sensitive data, if stolen, inadvertently shared, or exposed through a breach, can expose the individual and/or enterprise to various risks (e.g., identity theft or other crimes, embarrassment). As such, sensitive data included in or processed by an application needs to be protected, such as from unauthorized disclosure.

(2) It is with respect to these and other considerations that examples have been made. In addition, although relatively specific problems have been discussed, it should be understood that the examples should not be limited to solving the specific problems identified in the background.

### **SUMMARY**

(3) Examples described in this disclosure relate to systems and methods for protecting sensitive data. Data privacy is a concern when sensitive data, such as Personally Identifiable Information (PII), attorney-client privileged information, Protected Health Information (PHI), or other information that an individual or enterprise wants to keep from unauthorized disclosure is collected, stored, used, or otherwise processed in digital form. A challenge involving data privacy is to utilize data while protecting an individual and/or enterprise's privacy preferences/policies related to sensitive information. Sensitive data that is protected from unauthorized disclosure is herein referred to as privacy-relevant data. Examples of the present disclosure include a privacy-relevant data analyzer that analyzes source code, detects privacy-relevant data in the source code, and generates a report of instances of detected privacy-relevant data. In some examples, the privacy-relevant data analyzer scans through source code to detect annotations that denote if fields, records, or combinations thereof include privacy-relevant data. The privacy-relevant data analyzer further generates and provides a report of detected privacy issues associated with sensitive data included in source code so that the issues can be resolved to ensure that privacy is not breached.

(4) This Summary is provided to introduce a selection of concepts in a simplified form that are further described below in the Detailed Description. This Summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used to limit the scope of the claimed subject matter.

---

## **Description**

### **BRIEF DESCRIPTION OF THE DRAWINGS**

(1) The present disclosure is illustrated by way of example by the accompanying figures, in which like references indicate similar elements. Elements in the figures are illustrated for simplicity and clarity and have not necessarily been drawn to scale.

(2) FIG. 1 is a block diagram of a computing system environment for protecting privacy-relevant data according to an example;

(3) FIG. 2 is functional block diagram of an example computing system environment including a

privacy analyzer according to an example;

(4) FIG. 3 is an illustration of example source code including privacy-relevant data according to an example;

(5) FIG. 4A is an illustration of an example user interface as may be displayed during writing or editing source code according to an example;

(6) FIG. 4B is an illustration of an example notification as may be displayed in the example user interface based on an identified privacy-relevant data issue according to an example;

(7) FIG. 4C is an illustration of an example comment as may be added to the source code and displayed in the example user interface according to an example;

(8) FIG. 4D is an illustration of example sub-actions as may be displayed in the example user interface according to an example;

(9) FIG. 4E is an illustration of an example justification for suppressing a notification about a privacy-relevant data issue according to an example;

(10) FIG. 5 is a functional block diagram of an example computing system environment including a privacy analyzer according to an example;

(11) FIG. 6 is a flowchart depicting operations of a method for protecting sensitive data according to an example;

(12) FIG. 7 is a flowchart depicting operations of another method for protecting sensitive data according to an example;

(13) FIG. 8 is a flowchart depicting operations of another method for protecting sensitive data according to an example;

(14) FIG. 9 is a block diagram illustrating example physical components of a computing device with which aspects of the disclosure may be practiced; and

(15) FIGS. 10A and 10B are block diagrams of an example mobile computing device with which aspects of the present disclosure may be practiced.

#### DETAILED DESCRIPTION

(16) A software application's complexity can increase very quickly as the application is developed or updated. A codebase, for example, can span tens of thousands of lines of code, or more. It often becomes increasingly difficult to detect deficiencies such as privacy breaches, logic errors, security vulnerabilities, overlooked input possibilities, and even some typographical errors in the program or its underlying source code. Some errors, such as misspelled names or missing parentheses, may be caught by a compiler, but other deficiencies are not detected by compilers. For instance, the compiler will not catch a coding error that results in the application breaching privacy of sensitive data.

(17) As such, examples described in this disclosure relate to systems and methods for protecting privacy-relevant data from unauthorized disclosure in source code of an application. For instance, the present disclosure provides a plurality of technical features including: a privacy-relevant data analyzer that analyzes source code, detects privacy-relevant data in the source code, and generates a report of instances of detected privacy-relevant data. In some examples, the privacy-relevant data analyzer scans through source code to detect annotations that denote if fields, records, or combinations thereof include privacy-relevant data. The privacy-relevant data analyzer further generates and provides a report of detected privacy issues associated with sensitive data included in source code so that the issues can be resolved to ensure that privacy is not breached.

(18) According to examples, the privacy-relevant data analyzer improves data security by reducing coding errors that breach privacy of sensitive data. Moreover, annotations enrich data type descriptions with elements of a privacy taxonomy, allowing for scanning through vast amounts of source code data in memory efficiently, for example, with minimal runtime cost using machine-based reasoning. Thus, less memory space is required, processor load is reduced, and processing speed is increased, among other benefits.

(19) FIG. 1 is a block diagram illustrating an example computing system environment **100** in which

a privacy-relevant data analyzer **110** is implemented for protecting the privacy of sensitive data according to examples. The example computing system environment **100** as presented is a combination of interdependent components that interact to form an integrated whole. Components of the computing system environment **100** may be hardware components or software components (e.g., applications, application programming interfaces (APIs), modules, virtual machines, or runtime libraries) implemented on and/or executed by hardware components of the environment **100**. In one example, components of systems disclosed herein are implemented on a single processing device. The processing device provides an operating environment for software components to execute and utilize resources or facilities of such a system. An example of processing device(s) comprising such an operating environment is depicted in FIGS. **9**, **10A**, and **10B**. In another example, the components of systems disclosed herein are distributed across multiple processing devices. For instance, input may be entered on a client device and information may be processed on or accessed from other devices in a network, such as one or more remote cloud devices or web server devices.

(20) In FIG. **1**, the computing system environment **100** includes one or more general-purpose computing devices **102a**, **102b** (collectively, computing devices **102**). As can be appreciated, the scale and structure of the computing system environment **100** may vary and may include additional or fewer components than those described in FIG. **1**. As one example, one or more components included in the computing device(s) **102** may be incorporated into a service environment **106**.

(21) According to examples, a first computing device **102a** is used by a user, such as a developer. For instance, the user may be a developer who uses an interactive development environment (IDE) **112** operating on the first computing device **102a**. The IDE **112** is generally utilized to implement a programming environment that includes various tools to facilitate the development of source code **111** used in programs, applications, and other software solutions. A source code file (referred to herein generally as source code **111**) includes a set of instructions that is written in a programming language or a combination of programming languages. Source code **111** typically includes one or more statements, each statement typically including one or more expressions and/or entities. The expression and/or entity in the statement can be made up of multiple components.

(22) The IDE **112** typically enables developers/programmers (herein referred to generally as users) to write and edit source code **111**, see errors in code construction or syntax, automate repetitive tasks and the building of code assemblies, browse class structures, compile the source code into target code (e.g., JavaScript, low-level assembly language, binary object code, etc.), and the like. In some examples, the IDE **112** further provides code templates, macros, and other utilities; automatically creates classes, methods, and properties; supports code re-factoring; and supports tools for collaboration among development team members and project management, among other features.

(23) According to an example implementation, the computing system environment **100** further includes a privacy-relevant data analyzer **110**. According to examples, the IDE **112** utilizes the privacy-relevant data analyzer **110** to detect privacy-relevant data included in the source code **111**. For instance, the privacy-relevant data analyzer **110** evaluates source code **111**, detects privacy-relevant data included in the source code **111**, and generates a report of detected privacy-relevant data. In some examples, the report is provided to and read by the IDE **112**. For instance, the IDE **112** uses the report to notify the user of the detected privacy-relevant data items so that the user can review and resolve the items while the source code **111** is being written.

(24) According to another example implementation, the computing system environment **100** further includes a continuous integration tester **108** that analyzes source code **111**, for instance, as part of a code review process when integrating a source code part into a whole source code **111**. In some examples, and as depicted in FIG. **1**, the continuous integration tester **108** runs on a second computing device **102b**. According to examples, the continuous integration tester **108** utilizes the privacy-relevant data analyzer **110** to detect privacy-relevant data included in the source code **111**.

For instance, the privacy-relevant data analyzer **110** evaluates source code **111** for a current build, detects privacy-relevant data included in the source code **111**, and generates a report of detected privacy-relevant data issues. In some examples, the report is provided to and read by the continuous integration tester **108**. For instance, the continuous integration tester **108** uses the report to compare against a previous report from a previous build of the source code **111** and fails the current build when new unresolved sensitive data issues are detected. In some examples, the privacy-relevant data analyzer **110** is implemented as an application plugin to the IDE **112** and/or the continuous integration tester **108**. In other examples, the privacy-relevant data analyzer **110** is implemented as a stand-alone system. In other examples, the privacy-relevant data analyzer **110** is implemented as a component of the IDE **112**.

(25) According to examples, the computing devices **102** detect and/or collect input data from users or user devices. In some examples, the input data corresponds to user interaction with one or more software applications or services implemented by, or accessible to, the computing device **102**, such as the IDE **112** or and/or the continuous integration tester **108**. In other examples, the input data corresponds to automated interaction with the software applications or services, such as the automatic (e.g., non-manual) execution of scripts or sets of commands at scheduled times or in response to predetermined events. The user interaction or automated interaction may be related to performance of a user activity corresponding to a task, a project, a data request, or the like. The input data may include, for example, audio input, touch input, text-based input, gesture input, and/or image input. The input data may be detected/collected using one or more sensor components of computing devices **102**. Examples of sensors include microphones, touch-based sensors, geolocation sensors, accelerometers, optical/magnetic sensors, gyroscopes, keyboards, and pointing/selection tools. Examples of computing devices **102** include a personal computer (PC), mobile device (e.g., smartphone, tablet, laptop, personal digital assistant (PDA), wearable device (e.g., smart watch, smart eyewear, fitness tracker, smart clothing, body-mounted device, head-mounted display), gaming console or device, and an Internet of Things (IoT) device.

(26) According to an example implementation, the computing devices **102** provide the input data to the service environment **106** using a network **114**. Examples of the network **114** include a private area network (PAN), a local area network (LAN), a wide area network (WAN), and the like. Although the network **114** is depicted as a single network, it is contemplated that the network **114** may represent several networks of similar or varying types. The service environment **106** provides the computing devices **102** access to various computing services and resources (e.g., applications, devices, storage, processing power, networking, analytics, intelligence). As depicted in FIG. 1, the service environment **106** includes or provides access to one or more machine learning engines **116a**, **116b** (collectively, machine learning engines **116**). For instance, the machine learning engines **116** are implemented in some examples to learn rules for classifying data as privacy relevant. In some implementations, a first machine learning engine **116a** runs on the first computing device **102a**, providing a fast turnaround in the user's (e.g., developer's) workflow. In some implementations, a second machine learning engine **116b** runs on the second computing device **102b** (e.g., in a continuous integration/continuous deployment pipeline), providing slower but more thorough machine learning functionalities (e.g., corresponding to the software product being built).

(27) In some example implementations, the service environment **106** is implemented in a cloud-based or server-based environment using one or more computing devices, such as server devices (e.g., web servers, file servers, application servers, database servers), edge computing devices (e.g., routers, switches, firewalls, multiplexers), personal computers (PCs), virtual devices, and mobile devices. Alternatively, the service environment **106** may be implemented in an on-premises environment (e.g., a home or an office) using such computing devices. The computing devices comprise one or more sensor components, as discussed with respect to the computing device **102**. The service environment **106** may comprise numerous hardware and/or software components and may be subject to one or more distributed computing models/services (e.g., Infrastructure as a

Service (IaaS), Platform as a Service (PaaS), Software as a Service (SaaS), Functions as a Service (FaaS)).

(28) FIG. 2 is a functional block diagram of an example computing system environment **200** including the IDE **112** and the privacy-relevant data analyzer **110** for providing sensitive data protection in accordance with an embodiment of the present disclosure. The IDE **112** in the illustrated example includes a user interface (UI) **206**, which is typically implemented as a graphical user interface or “GUI”) that exposes development tools to a user, for example, to enable the user to generate source code **111** in a human-readable computer programming language (e.g., C#, Conceptual schema definition language (CSDL), Visual Basic, .NET programming language). The development tools, for example, may include a code editor **202**, a compiler **204**, an automation system **220**, a debugger **214**, and/or other tools.

(29) According to examples, source code **111** is created or modified by the code editor **202**, for instance, to develop or update an application. The source code **111** is compiled by the compiler **204** into executable code. In some examples, the compiler **204** is communicatively coupled to the UI **206** to expose errors to the user that may occur during compilation of the source code **111**. The code editor **202** is configured to enable source code **111** to be written and edited and will often include features to speed up input of source code, such as syntax highlighting, automated completion, and bracket matching functionality. In some examples, the code editor **202** is further configured to check syntax of the source code **111** on-the-fly as it is typed.

(30) The automation system **220** is configured to automate some of the tasks encountered when developing an application. The automation system **220** may include scripting or other automation tools to automate linking and compiling processes, for example, by performing scripted calls to the compiler **204**.

(31) The debugger **214** enables the programmer to observe run-time behavior of an application and locate logical and/or semantic errors in the source code **111**. For example, the debugger **214** allows the user to break/suspend execution of the application to examine the source code **111**, evaluate and edit variables in the application, view registers and instructions created from the source code **111**, and view the memory space used by the application. In some examples, the debugger **214** is configured to work with source code **111** at various stages in development.

(32) According to some examples, the IDE **112** interacts with external data sources that may be needed or helpful for the user to create the desired source code **111**. The external data sources, for example, may include files, references, data connections, libraries, and other items.

(33) According to some examples and as shown in FIG. 2, the IDE **112** interacts with the privacy-relevant data analyzer **110**. For instance, a request to analyze source code **111** that is being written or edited by the IDE **112** is received by the privacy-relevant data analyzer **110**. In some examples, the request corresponds to analyzing a portion of the source code **111**, such as an update or local change made to the source code **111**. In other examples, the request corresponds to analyzing at least a portion of the source code **111** that includes the update or local change. In response to receiving the request, the privacy-relevant data analyzer **110** evaluates the source code **111** for identifying privacy-relevant data included in the source code **111**, generates a report **222** based on the evaluation, and provides the report **222** to the IDE **112**. In some examples, the report **222** includes instructions for guiding an interaction with the user to resolve a detected privacy-relevant data issue in the source code **111**. In an example, the request for analysis and the response of the report **222** is performed as part of an interaction loop. For instance, the interaction loop is an ongoing loop that takes place while the source code **111** is being written or edited.

(34) In some example implementations, the privacy-relevant data analyzer **110** is implemented as a compiler that includes a variety of functional components including a lexical analyzer **208**, a parser **226**, a type checker **224**, data privacy analyzer **210**, and a report generator **218**. In some examples, one or more of the components are combined. In some examples, the components are implemented on a single computing device. In other examples, one or a combination of the components are

distributed across multiple computing devices.

(35) In some examples, the lexical analyzer **208** converts a stream of characters into a sequence of tokens which are defined, for example, by regular expressions. The parser **226** parses the token sequence to identify a syntactic structure of the application. A parse tree can be constructed to replace the linear structure of the token sequence by application of some formal grammar. For instance, a parse tree is a textual representation of source code **111** in a tree view according to specific language specification. The parse tree represents a given object model for source code **111** and is the basis on which privacy-relevance analysis is performed (e.g., by the data privacy analyzer **210**). In some implementations, additional semantic analysis is performed on the parse tree by performing type-checking processes (e.g., by the type checker **224**) to add semantic meaning to the parse tree. For instance, semantics is the logic in the source code **111** and all its constructs, which may include declarations of variables, classes, properties, objects, fields, methods, function calls and passing parameters to them, types of operands and operation results, and operator priorities.

(36) The data privacy analyzer **210** uses a set of inference rules **212** to determine whether a data field, record, or combination thereof in the source code **111** includes privacy-relevant data. In some examples, the data privacy analyzer **210** is configured to detect privacy-relevant data in classes, properties, and/or methods included in source code **111**.

(37) In some examples, privacy-relevant data is defined to include an annotation of one or more sets of privacy annotations. The one or more sets of privacy annotations may be input into source code **111** to denote whether a field, record, or combination thereof includes privacy-relevant data. In some examples, an annotation is manually added to a datatype in the source code **111**. In other examples, an annotation is machine learned and automatically added to the source code **111**. Inference rules **212** are used in this example to express various types of guidelines (and/or requirements) to which the source code **111** is desired to adhere and can comprise acceptable types of data the source code **111** can include and/or unacceptable types of data the source code **111** needs to protect (e.g., from disclosure). These can include performance guidelines for a program, design requirements, best practices, enterprise policies (or other types of corporate or company policies), legal guidelines, and regulatory and/or compliance guidelines, and various combinations thereof. In some examples, other types of guidelines may be utilized as may be needed to meet the requirements of a particular implementation. For example, law, regulations, rulings, edicts, or other type of imperatives to which strict adherence is needed can also be incorporated into the inference rules **212**. In some examples, the inference rules **212** encode the privacy guideline in a machine understandable format. For instance, the inference rules **212** are used to infer the data classification of a property from its components. As an example, a datatype is marked as containing privacy sensitive information by an inclusion of annotations associated with data classifications of properties of the datatype's components. If an inference rule states that the data classification of a particular property is a highest classification of its components, the datatype is classified as the particular privacy-relevant data.

(38) The report generator **218** generates a report **222** including a list of issues where instances of privacy-relevant data were determined and provides the report **222** to the IDE **112**. As mentioned above, in some examples, the report **222** further includes instructions for guiding an interaction with the user to resolve a detected privacy-relevant data issue in the source code **111**. For instance, the report **222** is utilized in a feedback loop to ensure that the output of the source code **111** complies with data privacy guidelines. In some examples, the IDE **112** maintains the list using a version control system. In some examples, when an issue is not handled or resolved, the IDE **112** is configured to mark the issue as an error such that the user's attention is directed to the issue.

(39) FIG. 3 shows a source code sample **300** (e.g., of source code **111**) that defines a datatype **304** in a formal language according to an example. In the illustrated example, the datatype **304** (e.g., “*TodoItemTDO*”) is a “class” type and includes a first data item **306a** (e.g., “*Id*”—a long integer)



and a second data item **306b** (e.g., “Name”—a string). In some examples and as shown in FIG. 3, annotations **302a**, **302b** (collectively, annotations **302**) are added to indicate which parts of the datatype **304** are privacy related. For instance, a first annotation **302a** (e.g., “EUPI” or “End User Pseudonymous Identifier”) is added to the first data item **306a** and a second annotation **302b** (e.g., “EUII” or “End User Identifiable Information”) is added to the second data item **306b**. The annotations **302** in this example are included in square brackets. In some examples, the annotations **302** are added to the source code **111** by the IDE **112** in response to manual user input of the annotations **302**. The annotations **302** are mapped to the annotation table **230**, which includes classifications of data that may be used to adhere to security, compliance, and privacy requirements and processes for collecting, storing, and using user personal information. The annotation table **230**, for example, represents a small set of annotative characters or symbols that are easily managed and manipulated as needed.

(40) FIG. 4A is a representation of an example UI **206** as may be displayed during writing or editing source code **111**. As depicted, the UI **206** displays a source code sample **400** (e.g., of source code **111**) including a method **402** of an application according to an example. For instance, the user may be writing or editing the method **402** using the code editor **202** of the IDE **112**. As depicted, the method **402** includes the datatype **304** defined in the source code sample **300** depicted in FIG. 3. For instance, the privacy-relevant data analyzer **110** may determine the method **402** includes an instance where privacy-relevant data may be exposed. In this example, the privacy-relevant data analyzer **110** evaluates the source code sample **400** and determines that the method **402** directly refers to (e.g., returns) the prior-defined datatype **304** that includes privacy-relevant data. According to an example, based on the annotations **302** included in the datatype **304**, the privacy-relevant data analyzer **110** identifies the datatype **304** is a violation of a guideline preventing leakage of the prior-defined datatype **304** and binds the datatype **304** as privacy-relevant (e.g., sensitive).

(41) In the other examples, the privacy-relevant data analyzer **110** evaluates the source code sample **400** and determines that the method **402** indirectly refers to the prior-defined datatype **304**. For instance, the method **402** may refer to a generic datatype, but may transform or otherwise process privacy-relevant data (e.g., the prior-defined datatype **304**) as a result of the execution of the method **402**. Accordingly, in some examples, the privacy-relevant data analyzer **110** is configured to use the set of inference rules **212** to infer whether a generic datatype is sensitive (e.g., whether a data field, record, or combination thereof the source code sample **400** includes privacy-relevant data).

(42) FIG. 4B depicts an example notification **404** as may be displayed in the example UI **206** based on an issue included in a report **222**. For instance, the UI **206** of the code editor **202** displays at least one line of the source code sample **400** that includes the determined instance of privacy-relevant data (e.g., datatype **304**). The UI **206** further displays the notification **404** in association with the privacy-relevant data instance. In some examples, the notification **404** points to or otherwise indicates the at least one line of the source code sample **400** that generated the issue and may include information about the issue and possible actions **406a**, **406b** (collectively, actions **406**) that may be taken to resolve or otherwise address or act on the issue. For instance, as can be seen in the example depicted in FIG. 4B, the endpoint is annotated with a “HttpGet” tag, and a comment describing the annotated element as an “endpoint” for an incoming HTTP request. Thus, the functionality exposed by the method **402**, is determined by inference to expose privacy sensitive information. The method **402**, by declaration, is exposing the privacy sensitive information through an external HTTP interface; therefore, the method **402** is a point in the source code **111** where privacy sensitive information is being exposed outside of the system and, thus, represents an endpoint that requires special privacy review. As an example, a first action **406a** may be selected to mark the method **402** as requiring special privacy review and a second action **406b** may be selected to suppress or configure the issue. For example, when an action **406** is selected, a corresponding

action to be performed is launched.

(43) FIG. 4C depicts an example comment **408** as may be displayed in the example UI **206** based on a selection of the first action **406a** included in the notification **404**. For instance, in response to receiving a selection to mark the method **402** as requiring special privacy review, the IDE **112** instructs the code editor **202** to insert the comment **408** into the source code **111**. As an example, the comment **408** may include information about the issue, a link for additional information or resources, etc.

(44) FIG. 4D depicts example sub-actions **410a**, **410b** (collectively, sub-actions **410**) as may be displayed in the example UI **206** based on a selection of the second action **406b** included in the notification **404** depicted in FIG. 4B. For instance, in response to receiving a selection to suppress or configure the issue, one or more sub-actions **410** are presented in the UI **206**, such as a first sub-action **410a** that may be selected to suppress the issue in the source code **111** and a second sub-action **410b** that may be selected to configure the severity for the privacy-relevant data analyzer **110**. As shown, the user selects the first sub-action **410a**, and with reference now to FIG. 4E, an example justification **412** is depicted as may be added to the source code **111** and displayed in the example UI **206** based on a selection of the first sub-action **410a**. For instance, when a selection is made to suppress the issue, an annotation or assembly decorator is generated. In some examples, the user is prompted to enter a justification **412** for the selection to suppress the notification, which is included in the annotation and displayed in the UI **206**. In other examples, the justification **412** is included in a same or separate file, stored, and is tracked (e.g., using a source control system). According to examples, the IDE **112** tracks issues that need to be reviewed. In some examples, when an issue is suppressed and not resolved, the IDE **112** tracks the issue as an error in the source code **111**.

(45) FIG. 5 is a functional block diagram of an example computing system environment **500** including the privacy-relevant data analyzer **110** and the continuous integration tester **108** for providing sensitive data privacy protection in accordance with another embodiment of the present disclosure. For instance, the example computing system environment **200** described above with reference to FIG. 2 may operate to perform operations in a build pipeline **502**, where the example computing system environment **500** with reference to FIG. 5 may operate to perform operations in a continuous integration pipeline **504**. In the continuous integration pipeline **504**, for example, the continuous integration tester **108** takes artifacts generated in the build pipeline **502** (e.g., the source code **111** and a previous report **222a** from a previous build) and tests the source code **111** before sending passing source code **111** (e.g., source code that passes the test) to an execution environment **508** (e.g., a substrate or other execution environment). In some examples, when source code **111** fails the test, the continuous integration tester **108** fails the current build. For instance, an application typically evolves in stages, where different stages of the source code **111** are ongoingly deployed (e.g., later a same day, a next day, a next week, a next year). Accordingly, each new version of the source code **111** that is being deployed is tested by the continuous integration tester **108**. As such, in some examples, the continuous integration tester **108** requests source code **111** that is being written or edited by a user. Additionally, the continuous integration tester **108** requests the previous report **222a** from the previous build or deployment of the source code **111**. For instance, the continuous integration tester **108** uses the previous report **222a** to find differences between old and new source code **111**. In some examples, the continuous integration tester **108** runs compilation and static analyses of source code **111** as a batch job.

(46) In some examples, a request to analyze source code **111** in the continuous integration pipeline **504** is received by the privacy-relevant data analyzer **110** (from the continuous integration tester **108**) in association with a current build. In some examples, the request corresponds to analyzing a portion of the source code **111**, such as an update or local change made to the source code **111** in the current build. In other examples, the request corresponds to analyzing at least a portion of the source code **111** that includes the update or local change. In response to receiving the request, the

privacy-relevant data analyzer **110** evaluates the source code **111** for identifying privacy-relevant data included in the source code **111**, generates a report (e.g., a current report **222b**) based on the evaluation, and provides the current report **222b** to the continuous integration tester **108**. In some examples, the continuous integration tester **108** compares the current report **222b** to the previous report **222a** for determining whether there are new or unresolved sensitive data issues in the current report **222b**. In some examples, when the continuous integration tester **108** detects that a privacy-relevant data issue is not resolved, the continuous integration tester **108** fails the current build so that the privacy-relevant data issue is not promoted into any kind of further integration. For instance, if a privacy-relevant data issue is not explicitly resolved, a link between the continuous integration pipeline **504** and continuous deployment pipeline **506** is broken and the source code **111** is not promoted into production status.

(47) FIG. **6** is a flowchart depicting a method **600** for protecting sensitive data according to an example. For instance, one or more operations of the method **600** are performed by the IDE **112** as part of protecting sensitive data. With reference now to FIG. **6**, the method **600** starts when source code **111** being written or edited is received by the IDE **112**. For instance, a user may be writing or editing source code **111** of an application during in a build pipeline **502** of a development process. In some examples, the source code **111** is a first version of source code **111**. In other examples, the source code **111** is an iterative version of source code **111**.

(48) At operation **604**, the IDE **112** requests data privacy analysis of at least a portion of the source code **111** including an update or local change made to the source code **111** by the user. For example, the IDE **112** requests the privacy-relevant data analyzer **110** to evaluate the source code **111** for determining whether the update or local change exposes sensitive data. In some examples, the privacy-relevant data analyzer **110** performs one or more operations of a method **800** depicted in FIG. **8** and described below for protecting sensitive data according to an example.

(49) At operation **606**, a report **222** is received from the privacy-relevant data analyzer **110** in association with the evaluation. At decision operation **608**, a determination is made as to whether the report **222** includes any privacy-relevant data issues corresponding to one or more instances of detected privacy-relevant data in the source code update or local change. When a privacy-relevant data issue is included in the report **222**, a review list is generated at operation **610** to track each privacy-relevant data issue, such as whether the issue has been reviewed, resolved, suppressed, etc.

(50) At operation **612**, a notification **404** is generated for display in a UI **206** in association with the privacy-relevant data instance. In some examples, the notification **404** points to or otherwise indicates a datatype **304** that generated the issue and may include information about the issue and possible actions **406** that may be taken to resolve or otherwise address the issue.

(51) At decision operation **614**, a determination is made as to whether the issue has been resolved. For instance, the issue is marked as resolved at operation **616** when an action **406** selected by the user that causes the issue to be reviewed is explicitly resolved. Alternatively, when the issue is not resolved, an error is generated at operation **618** in the source code **111**. In some examples, the method **600** returns to operation **612**, where a notification of the issue is generated and displayed to the user. For instance, the user may then be able to resolve or otherwise address the issue.

(52) FIG. **7** is a flowchart depicting a method **700** for protecting sensitive data according to an example. For instance, one or more operations of the method **700** are performed by the continuous integration tester **108** as part of protecting sensitive data. With reference now to FIG. **7**, the method **700** starts at operation **702** when source code **111** of a current build of an application is received by the continuous integration tester **108** as part of a continuous integration pipeline **504**. For instance, the current build includes an update made to the source code **111** of the application from a previous build.

(53) At operation **704**, the continuous integration tester **108** requests and receives a previous privacy report **222a** of the previous build of the source code **111**.

(54) At operation **706**, the continuous integration tester **108** requests data privacy analysis of at

least a portion of the source code **111** including the update or local change made to the source code **111** by the user. For example, the continuous integration tester **108** requests the privacy-relevant data analyzer **110** to evaluate the source code **111** for determining whether the update or local change exposes sensitive data. In some examples, the privacy-relevant data analyzer **110** performs one or more operations of a method **800** depicted in FIG. **8** and described below for protecting sensitive data according to an example.

(55) At operation **708**, a current report **222b** is received from the privacy-relevant data analyzer **110** in association with the evaluation. At decision operation **710**, the privacy-relevant data analyzer **110** compares the current report **222b** and the previous report **222a** and determines whether the current report **222b** includes any unresolved or new privacy-relevant data issues. For example, the unresolved or new privacy-relevant data issues may correspond to one or more instances of detected privacy-relevant data in the source code update or local change. In some examples, when a new or unresolved privacy-relevant data issue is determined, the privacy-relevant data analyzer **110** fails the current build at operation **712**. In some examples, the privacy-relevant data analyzer **110** generates a justification for the failed build. The justification may include an indication of the unresolved or new privacy-relevant data issue. In other examples, when no new or unresolved privacy-relevant data issues are determined at decision operation **710**, the privacy-relevant data analyzer **110** passes the current build at operation **714**.

(56) FIG. **8** is a flowchart depicting a method **700** for protecting sensitive data according to an example. For instance, one or more operations of the method **800** are performed by the privacy-relevant data analyzer **110**. In some examples, the method **800** is performed during a build pipeline **502** of development of an application. In other examples, the method **800** is performed during a continuous integration pipeline **504** of the development process.

(57) At operation **802**, at least a portion of source code **111** that includes an update or local change is received in a request for privacy analysis. In response to receiving the request, the privacy-relevant data analyzer **110** analyzes the source code **111** at operation **804** for identifying privacy-relevant data included in the source code **111**. In some examples, the privacy-relevant data analyzer **110** generates a parse tree representing the source code **111** and performs an analysis of the parse tree for determining whether the source code **111** includes privacy-relevant data. For instance, the privacy-relevant data analyzer **110** performs a traversal of the parse tree to determine whether the source code **111** includes privacy-relevant data classifications.

(58) At decision operation **806**, the privacy-relevant data analyzer **110** determines whether the source code **111** includes privacy-relevant data. In some examples, the determination is based on whether the parse tree includes a direct reference to a datatype **304** including an annotation **302** mapped to annotation table **230** including classifications of data and their associated annotations **302** that may be used to adhere to security, compliance, and privacy requirements and processes for collecting, storing, and using user personal information. In other examples, the determination is made based on inferred whether a datatype **304** referred to in a method **402** includes privacy-relevant data. The inference, for example, may be made based on a set of inference rules **212**. As an example, a parameter to a certain logging function is defined to only accept data with certain classifications, (e.g., “System Metadata”, which is not sensitive). For a variable/function which is not explicitly annotated with a classification, the variable will receive the sum of all classifications it had as input (i.e., a worst-case interpretation). Thus, if any data is sent (e.g., transitively) classified as private in that parameter, the privacy-relevant data analyzer **110** determines this as breaking the rule and triggering an error.

(59) When a determination is made that the source code **111** includes privacy-relevant data, the privacy-relevant data analyzer **110** generates a report **222** at operation **808**. According to examples, the report **222** includes an information about one or more privacy-relevant issues corresponding to one or more instances of privacy-relevant data included in the source code as determined by the privacy-relevant data analyzer **110**.

(60) At operation **810**, the privacy-relevant data analyzer **110** provides the report **222** to a requestor. For example, when the method **800** is performed during the build pipeline **502** of the application, the privacy-relevant data analyzer **110** provides the report **222** to the IDE **112**. Or, when the method **800** is performed during the continuous integration pipeline **504** of the development process of the application, the privacy-relevant data analyzer **110** provides the report **222** to the continuous integration tester **108**.

(61) FIGS. **9**, **10A**, and **10B** and the associated descriptions provide a discussion of a variety of operating environments in which examples of the disclosure may be practiced. However, the devices and systems illustrated and discussed with respect to FIGS. **9**, **10A**, and **10B** are for purposes of example and illustration, a vast number of computing device configurations that may be utilized for practicing aspects of the disclosure, described herein.

(62) FIG. **9** is a block diagram illustrating physical components (e.g., hardware) of a computing device **900** with which examples of the present disclosure may be practiced. The computing device components described below may be suitable for one or more of the components of the system **100** described above. In a basic configuration, the computing device **900** includes at least one processing unit **902** and a system memory **904**. Depending on the configuration and type of computing device **900**, the system memory **904** may comprise volatile storage (e.g., random access memory), non-volatile storage (e.g., read-only memory), flash memory, or any combination of such memories. The system memory **904** may include an operating system **905** and one or more program modules **906** suitable for running software applications **950**, such as the IDE **112**, privacy-relevant data analyzer **110**, continuous integration tester **108**, and other applications.

(63) The operating system **905** may be suitable for controlling the operation of the computing device **900**. Furthermore, aspects of the disclosure may be practiced in conjunction with a graphics library, other operating systems, or any other application program and is not limited to any particular application or system. This basic configuration is illustrated in FIG. **9** by those components within a dashed line **908**. The computing device **900** may have additional features or functionality. For example, the computing device **900** may also include additional data storage devices (removable and/or non-removable) such as, for example, magnetic disks, optical disks, or tape. Such additional storage is illustrated in FIG. **9** by a removable storage device **909** and a non-removable storage device **910**.

(64) As stated above, a number of program modules and data files may be stored in the system memory **904**. While executing on the processing unit **902**, the program modules **906** may perform processes including one or more of the stages of methods **600**, **700**, and/or **800** illustrated in FIG. **5**, FIG. **6**, and/or FIG. **8**, respectively. Other program modules that may be used in accordance with examples of the present disclosure and may include applications such as electronic mail and contacts applications, word processing applications, spreadsheet applications, database applications, slide presentation applications, drawing or computer-aided application programs, etc.

(65) Furthermore, examples of the disclosure may be practiced in an electrical circuit comprising discrete electronic elements, packaged or integrated electronic chips containing logic gates, a circuit utilizing a microprocessor, or on a single chip containing electronic elements or microprocessors. For example, examples of the disclosure may be practiced via a system-on-a-chip (SOC) where each or many of the components illustrated in FIG. **9** may be integrated onto a single integrated circuit. Such an SOC device may include one or more processing units, graphics units, communications units, system virtualization units and various application functionality all of which are integrated (or “burned”) onto the chip substrate as a single integrated circuit. When operating via an SOC, the functionality, described herein, with respect to providing spatial-textual clustering-based predictive recognition of text in a video may be operated via application-specific logic integrated with other components of the computing device **900** on the single integrated circuit (chip). Examples of the present disclosure may also be practiced using other technologies capable of performing logical operations such as, for example, AND, OR, and NOT, including mechanical,

optical, fluidic, and quantum technologies.

(66) The computing device **900** may also have one or more input device(s) **912** such as a keyboard, a mouse, a pen, a sound input device, a touch input device, a camera, etc. The output device(s) **914** such as a display, speakers, a printer, etc. may also be included. The aforementioned devices are examples and others may be used. The computing device **900** may include one or more communication connections **916** allowing communications with other computing devices **918**. Examples of suitable communication connections **916** include RF transmitter, receiver, and/or transceiver circuitry; universal serial bus (USB), parallel, and/or serial ports.

(67) The term computer readable media as used herein includes volatile and nonvolatile, removable and non-removable media implemented in any method or technology for storage of information, such as computer readable instructions, data structures, or program modules. The system memory **904**, the removable storage device **909**, and the non-removable storage device **910** are all computer readable media examples (e.g., memory storage.) Computer readable media include random access memory (RAM), read-only memory (ROM), electrically erasable programmable ROM (EEPROM), flash memory or other memory technology, CD-ROM, digital versatile disks (DVD) or other optical storage, magnetic cassettes, magnetic tape, magnetic disk storage or other magnetic storage devices, or any other article of manufacture which can be used to store information and which can be accessed by the computing device **900**. Any such computer readable media may be part of the computing device **900**. Computer readable media does not include a carrier wave or other propagated data signal.

(68) Communication media may be embodied by computer readable instructions, data structures, program modules, or other data in a modulated data signal, such as a carrier wave or other transport mechanism, and includes any information delivery media. The term “modulated data signal” may describe a signal that has one or more characteristics set or changed in such a manner as to encode information in the signal. By way of example, communication media may include wired media such as a wired network or direct-wired connection, and wireless media such as acoustic, RF, infrared, and other wireless media.

(69) FIGS. **10A** and **10B** illustrate a mobile computing device **1000**, for example, a mobile telephone, a smart phone, a tablet personal computer, a laptop computer, and the like, with which aspects of the disclosure may be practiced. With reference to FIG. **10A**, an example of a mobile computing device **1000** for implementing at least some aspects of the present technology is illustrated. In a basic configuration, the mobile computing device **1000** is a handheld computer having both input elements and output elements. The mobile computing device **1000** typically includes a display **1005** and one or more input buttons **1010** that allow the user to enter information into the mobile computing device **1000**. The display **1005** of the mobile computing device **1000** may also function as an input device (e.g., a touch screen display). If included, an optional side input element **1015** allows further user input. The side input element **1015** may be a rotary switch, a button, or any other type of manual input element. In alternative examples, mobile computing device **1000** may incorporate more or less input elements. For example, the display **1005** may not be a touch screen in some examples. In alternative examples, the mobile computing device **1000** is a portable phone system, such as a cellular phone. The mobile computing device **1000** may also include an optional keypad **1035**. Optional keypad **1035** may be a physical keypad or a “soft” keypad generated on the touch screen display. In various aspects, the output elements include the display **1005** for showing a graphical user interface (GUI), a visual indicator **1020** (e.g., a light emitting diode), and/or an audio transducer **1025** (e.g., a speaker). In some examples, the mobile computing device **1000** incorporates a vibration transducer for providing the user with tactile feedback. In yet another example, the mobile computing device **1000** incorporates input and/or output ports, such as an audio input (e.g., a microphone jack), an audio output (e.g., a headphone jack), and a video output (e.g., a HDMI port) for sending signals to or receiving signals from an external device.

(70) FIG. 10B is a block diagram illustrating the architecture of one example of a mobile computing device. That is, the mobile computing device **1000** can incorporate a system (e.g., an architecture) **1002** to implement some examples. In one example, the system **1002** is implemented as a “smart phone” capable of running one or more applications (e.g., videoconference or virtual meeting application, browser, e-mail, calendaring, contact managers, messaging clients, games, and media clients/players). In some examples, the system **1002** is integrated as a computing device, such as an integrated personal digital assistant (PDA) and wireless phone.

(71) One or more application programs **1050** (e.g., one or more of the components of system **100**) may be loaded into the memory **1062** and run on or in association with the operating system **1064**, such as the IDE **112**, privacy-relevant data analyzer **110**, continuous integration tester **108**, and/or other applications. Other examples of the application programs **1050** include videoconference or virtual meeting programs, phone dialer programs, e-mail programs, personal information management (PIM) programs, word processing programs, spreadsheet programs, Internet browser programs, messaging programs, and so forth. The system **1002** also includes a non-volatile storage area **1068** within the memory **1062**. The non-volatile storage area **1068** may be used to store persistent information that should not be lost if the system **1002** is powered down. The application programs **1050** may use and store information in the non-volatile storage area **1068**, such as e-mail or other messages used by an e-mail application, and the like. A synchronization application (not shown) also resides on the system **1002** and is programmed to interact with a corresponding synchronization application resident on a host computer to keep the information stored in the non-volatile storage area **1068** synchronized with corresponding information stored at a remote device or server. As should be appreciated, other applications may be loaded into the memory **1062** and run on the mobile computing device **1000**.

(72) The system **1002** has a power supply **1070**, which may be implemented as one or more batteries. The power supply **1070** might further include an external power source, such as an AC adapter or a powered docking cradle that supplements or recharges the batteries.

(73) The system **1002** may also include a radio **1072** that performs the function of transmitting and receiving radio frequency (RF) communications. The radio **1072** facilitates wireless connectivity between the system **1002** and the “outside world,” via a communications carrier or service provider. Transmissions to and from the radio **1072** are conducted under control of the operating system **1064**. In other words, communications received by the radio **1072** may be disseminated to the application programs **1050** via the operating system **1064**, and vice versa.

(74) The visual indicator **1020** (e.g., light emitting diode (LED)) may be used to provide visual notifications and/or an audio interface **1074** may be used for producing audible notifications via the audio transducer **1025**. In the illustrated example, the visual indicator **1020** is a light emitting diode (LED) and the audio transducer **1025** is a speaker. These devices may be directly coupled to the power supply **1070** so that when activated, they remain on for a duration dictated by the notification mechanism even though the processor **1060** and other components might shut down for conserving battery power. The LED may be programmed to remain on indefinitely until the user takes action to indicate the powered-on status of the device. The audio interface **1074** is used to provide audible signals to and receive audible signals from the user. For example, in addition to being coupled to the audio transducer **1025**, the audio interface **1074** may also be coupled to a microphone to receive audible input, such as to facilitate a telephone conversation. The system **1002** may further include a video interface **1076** that enables an operation of a peripheral device port **1030** (e.g., an on-board camera) to record still images, video stream, and the like.

(75) A mobile computing device **1000** implementing the system **1002** may have additional features or functionality. For example, the mobile computing device **1000** may also include additional data storage devices (removable and/or non-removable) such as, magnetic disks, optical disks, or tape. Such additional storage is illustrated in FIG. 10B by the non-volatile storage area **1068**.

(76) Data/information generated or captured by the mobile computing device **1000** and stored via

the system **1002** may be stored locally on the mobile computing device **1000**, as described above, or the data may be stored on any number of storage media that may be accessed by the device via the radio **1072** or via a wired connection between the mobile computing device **1000** and a separate computing device associated with the mobile computing device **1000**, for example, a server computer in a distributed computing network, such as the Internet. As should be appreciated such data/information may be accessed via the mobile computing device **1000** via the radio **1072** or via a distributed computing network. Similarly, such data/information may be readily transferred between computing devices for storage and use according to well-known data/information transfer and storage means, including electronic mail and collaborative data/information sharing systems.

(77) Examples include a computer-implemented method for protecting sensitive data, comprising: receiving a request from a requestor to analyze source code of an application during development of the application; analyzing the source code for privacy relevant data based on a set of inference rules; determining an instance of privacy related data corresponding to an annotation included in the source code; generating a report including an indication of a privacy issue corresponding to the instance of privacy related data; and providing the report to the requestor for providing a notification of the privacy issue.

(78) Examples include a system, the system comprising at least one processor; and memory storing instructions that, when executed by the at least one processor cause the system to: receive an update to source code of an application during development of the application; analyze the source code for privacy relevant data based on a set of inference rules; determine one or more instances of privacy related data corresponding to one or more annotations included in the source code; generate a first report including an indication of one or more privacy issues corresponding to the one or more instances of privacy related data; and generate a notification for each of the one or more privacy issues in the first report that each privacy issue should be reviewed.

(79) Examples include a system comprising: at least one processor; and memory storing instructions that, when executed by the at least one processor cause the system to: receive an update to source code of an application during development of the application; analyze the source code for privacy relevant data based on a set of inference rules; determine one or more instances of privacy related data corresponding to one or more annotations included in the source code; generate a first report including an indication of one or more privacy issues corresponding to the one or more instances of privacy related data; and generate a notification for each of the one or more privacy issues in the first report that each privacy issue should be reviewed.

(80) The methods, modules, and components depicted herein are merely examples. Alternatively, or in addition, the functionality described herein can be performed, at least in part, by one or more hardware logic components. For example, illustrative types of hardware logic components that can be used include Field-Programmable Gate Arrays (FPGAs), Application-Specific Integrated Circuits (ASICs), Application-Specific Standard Products (ASSPs), System-on-a-Chip systems (SOCs), Complex Programmable Logic Devices (CPLDs), etc. In an abstract, but still definite sense, any arrangement of components to achieve the same functionality is effectively “associated” such that the desired functionality is achieved. Hence, any two components herein combined to achieve a particular functionality can be seen as “associated with” each other such that the desired functionality is achieved, irrespective of architectures or inter-medial components. Likewise, any two components so associated can also be viewed as being “operably connected,” or “coupled,” to each other to achieve the desired functionality. Merely because a component, which may be an apparatus, a structure, a system, or any other implementation of a functionality, is described herein as being coupled to another component does not mean that the components are necessarily separate components. As an example, a component A described as being coupled to another component B may be a sub-component of the component B, the component B may be a sub-component of the component A, or components A and B may be a combined sub-component of another component C.

(81) Furthermore, boundaries between the functionality of the above described operations are



merely illustrative. The functionality of multiple operations may be combined into a single operation, and/or the functionality of a single operation may be distributed in additional operations. Moreover, alternative embodiments may include multiple instances of a particular operation, and the order of operations may be altered in various other embodiments.

(82) Although the disclosure provides specific examples, various modifications and changes can be made without departing from the scope of the disclosure as set forth in the claims below.

Accordingly, the specification and figures are to be regarded in an illustrative rather than a restrictive sense, and all such modifications are intended to be included within the scope of the present disclosure. Any benefits, advantages, or solutions to problems that are described herein with regard to a specific example are not intended to be construed as a critical, required, or essential feature or element of any or all the claims.

(83) Furthermore, the terms “a” or “an,” as used herein, are defined as one or more than one. Also, the use of introductory phrases such as “at least one” and “one or more” in the claims should not be construed to imply that the introduction of another claim element by the indefinite articles “a” or “an” limits any particular claim containing such introduced claim element to containing only one such element, even when the same claim includes the introductory phrases “one or more” or “at least one” and indefinite articles such as “a” or “an.” The same holds true for the use of definite articles.

(84) Unless stated otherwise, terms such as “first” and “second” are used to arbitrarily distinguish between the elements such terms describe. Thus, these terms are not necessarily intended to indicate temporal or other prioritization of such elements.

## Claims

1. A computer-implemented method for protecting sensitive data, comprising: receiving a request from a requestor to analyze source code of an application during development of the application; analyzing the source code for privacy relevant data based on a set of inference rules; determining an instance of privacy related data corresponding to an annotation included in the source code, wherein the annotation is mapped to an annotation table comprising a mapping of the annotation to a privacy relevant data classification, and wherein the instance of privacy related data is determined based on the privacy relevant data classification; generating a report including an indication of a privacy issue corresponding to the instance of privacy related data; and providing the report to the requestor for providing a notification of the privacy issue.
2. The method of claim 1, wherein providing the report to the requestor comprises: including, in the report, information corresponding to the instance of privacy related data; and providing the report to an interactive development environment (IDE) during a build pipeline of the development of the application.
3. The method of claim 2, wherein providing the report comprises including instructions for generating a notification notifying a user of the privacy issue.
4. The method of claim 1, wherein providing the report to the requestor comprises: including, in the report, information corresponding to a new or unresolved privacy related data; and providing the report to a continuous integration tester during a continuous integration pipeline of the development of the application for failing a current build based on the new or unresolved privacy related data.
5. The method of claim 1, wherein analyzing the source code for privacy relevant data comprises using one or more inference rules defining privacy-relevant data.
6. The method of claim 1, wherein analyzing the source code for privacy relevant data comprises determining whether the source code includes a class, property, or method that refers to a datatype including the annotation.
7. A system comprising: at least one processor; and memory storing instructions that, when

executed by the at least one processor cause the system to: receive a request from a requestor to analyze source code of an application during development of the application; analyze the source code for privacy relevant data based on a set of inference rules; determine an instance of privacy related data corresponding to an annotation included in the source code, wherein the annotation is mapped to an annotation table comprising a mapping of the annotation to a privacy relevant data classification, and wherein the instance of privacy related data is determined based on the privacy relevant data classification; generate a report including an indication of a privacy issue corresponding to the instance of privacy related data; and provide the report to the requestor for providing a notification of the privacy issue.

8. The system of claim 7, wherein the instructions cause the system to: include, in the report, information corresponding to the instance of privacy related data; and provide the report to an interactive development environment (IDE) during a build pipeline of the development of the application.

9. The system of claim 8, wherein the instructions cause the system to include, in the report, instructions for generating a notification notifying a user of the privacy issue.

10. The system of claim 7, wherein the instructions cause the system to: include, in the report, information corresponding to a new or unresolved privacy related data; and provide the report to a continuous integration tester during a continuous integration pipeline of the development of the application for failing a current build based on the new or unresolved privacy related data.

11. The system of claim 7, wherein in analyzing the source code, the instructions cause the system to use one or more inference rules defining privacy-relevant data to determine a class, property, or method refers to a datatype including the annotation.

12. The system of claim 11, wherein the annotation is manually added to the datatype.

13. The system of claim 11, wherein the annotation is machine learned and automatically added to the datatype.

14. A system comprising: at least one processor; and memory storing instructions that, when executed by the at least one processor cause the system to: receive an update to source code of an application during development of the application; analyze the source code for privacy relevant data based on a set of inference rules; determine one or more instances of privacy related data corresponding to one or more annotations included in the source code, wherein the annotation is mapped to an annotation table comprising a mapping of the one or more annotations to a privacy relevant data classification, and wherein the instance of privacy related data is determined based on the privacy relevant data classification; generate a first report including an indication of one or more privacy issues corresponding to the one or more instances of privacy related data; and generate a notification for each of the one or more privacy issues in the first report that each privacy issue should be reviewed.

15. The system of claim 14, wherein in analyzing the source code, the instructions cause the system to use one or more inference rules defining privacy-relevant data to determine a class, property, or method refers to a datatype including an annotation of the one or more annotations.

16. The system of claim 15, wherein the annotation of the one or more annotations is manually added to the datatype.

17. The system of claim 15, wherein the annotation of the one or more annotations is machine learned and automatically added to the datatype.

18. The system of claim 14, wherein the instructions cause the system to: determine at least one of the one or more privacy issues in the first report is unresolved; and generate an error corresponding to the at least one unresolved privacy issue.

19. The system of claim 14, wherein the instructions cause the system to: test the source code before sending the source code to an execution environment; analyze the source code for new or unresolved privacy relevant data based on a set of inference rules; determine one or more instances of new or unresolved privacy related data corresponding to one or more annotations included in the

source code; generate a second report including an indication of one or more privacy issues corresponding to the one or more instances of new or unresolved privacy related data; and fail the source code and prevent the source code from being sent to the execution environment.

20. The system of claim 19, wherein in determining one or more instances of new or unresolved privacy related data, the instructions cause the system to compare the second report to the first report for identifying the one or more instances of new or unresolved privacy related data.

---