

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12393729
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Anderson; Erik Daniel et al.

Privacy preserving user personalization using noisy ranking

Abstract

Disclosed is a system that tracks website usage without compromising user privacy. The system aggregates website usage data of multiple users across multiple websites. Website usage data is aggregated in a way that preserves each individual user's privacy. Specifically, information relevant to a particular user may be obtained from the aggregated information without exposing what was actually collected from that user. In some configurations, user-specific website usage data is aggregated using trusted execution environment computing hardware. This ensures that privacy is preserved while user-specific data is transferred to and processed by the system. The trusted execution environment applies differential privacy techniques to ensure that use of the aggregated information does not reveal actual information about a user's website usage history. In this way, privacy is maintained while still enabling many of the scenarios that would otherwise rely on third-party cookies.

Inventors: Anderson; Erik Daniel (Redmond, WA), Pfeiffer, III; Joseph J. (Seattle, WA), Charles; Denis Xavier (Redmond, WA), Rebrikov; Aleksandr (Bellevue, WA), Mooring; John Robert (Bellevue, WA), Maslen; Brandon J. (Redmond, WA), Gilton; Davis Leland (Seattle, WA), Yekhanin; Sergey (Redmond, WA), Gopi; Sivakanth (Skokie, IL)

Applicant: MICROSOFT TECHNOLOGY LICENSING, LLC (Redmond, WA)

Family ID: 1000008764769

Assignee: MICROSOFT TECHNOLOGY LICENSING, LLC (Redmond, WA)

Appl. No.: 18/301943

Filed: April 17, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240346178 A1	Oct. 17, 2024

Publication Classification

Int. Cl.: G06F21/62 (20130101); G06F21/60 (20130101)

U.S. Cl.:

CPC G06F21/6254 (20130101); G06F21/602 (20130101);

Field of Classification Search

CPC: G06F (21/6254); G06F (21/602)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
9002975	12/2014	Shenfield et al.	N/A	N/A
10129211	12/2017	Heath	N/A	G06Q 30/02
10636056	12/2019	Chakraborty et al.	N/A	N/A
2002/0007453	12/2001	Nemovicher	713/170	H04L 63/105
2013/0061035	12/2012	Hook	713/150	G06F 21/62
2020/0089917	12/2019	Skourtis	N/A	G06F 21/76
2021/0217053	12/2020	Yu	N/A	N/A
2023/0043318	12/2022	Selman	N/A	G06F 21/6209

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
114298734	12/2021	CN	N/A
2008539508	12/2007	JP	N/A

OTHER PUBLICATIONS

Anders, Kelda, "Introduction", Retrieved From: <https://github.com/WICG/privacy-preserving-ads/blob/main/NoisyRanking.md>, Apr. 26, 2022, 25 Pages. cited by applicant

Primary Examiner: Sall; El Hadji M

Attorney, Agent or Firm: Newport IP, LLC

Background/Summary

BACKGROUND

(1) A cookie is a small piece of data stored on a user's device by a website. Cookies are created by web servers and sent to the user's browser where they store information such as preferences, login details, or items in a shopping cart. First party cookies originate from a website the user chose to visit. Third party cookies originate from a different domain than the website that the user chose to visit. For example, a news website from one domain may embed a weather forecast website from a different domain. Both the news website and the weather forecast website may download cookies

onto the user's device. Cookies from the news website are considered first party because the user chose to navigate to the news website. Cookies from the weather website are considered third-party because they originated from a domain that the user did not explicitly choose to visit.

(2) Third-party cookies are often used for tracking purposes, such as tracking a user's online behavior or serving targeted content. For example, a third-party analytics service may use third-party cookies to record which pages of a website a user has visited and how long they spent on each page. The website owner may use this information to improve the user experience or optimize the performance of the website.

(3) Third-party cookies are also used to track a user across multiple domains, allowing the third-party to build a detailed profile of the user's online behavior and interests. This data can be used to improve the user's experience by providing content that is more relevant to the user's interests, and can also help to generate revenue for website owners. For example, a news website, an email website, and a social media website may all embed a web component from a third-party search provider. As the user browses these websites the third-party web component may store user interactions in a third-party cookie. The third-party cookie will also be accessible when the user visits the search provider's website directly. This gives the third-party visibility into the user's behavior while they were reading the news, checking their email, and browsing their social media feed.

(4) One concern with third-party cookies is invasion of a user's privacy, as they allow the third-party to track a user's online activity without the user's explicit knowledge or consent. This is particularly worrisome when a third-party has an expansive reach, allowing it to track a significant percentage of a user's web browsing activity. In response to these privacy concerns, many web browsers now include options for blocking or restricting the use of third-party cookies.

Unfortunately, while blocking third-party cookies protects a user's privacy, it also results in users receiving less relevant content suggestions, websites no longer having the insight provided by third-party analytics services, and the loss of other benefits enabled by third-party cookies.

(5) It is with respect to these and other considerations that the disclosure made herein is presented.

SUMMARY

(6) Disclosed is a system that delivers content without compromising user privacy. The system aggregates website usage data of multiple users across multiple websites. Website usage data is aggregated in a way that preserves each individual user's privacy. Specifically, information relevant to a particular user may be obtained from the aggregated information without exposing the granular identifying information that was collected from that user. In some configurations, user-specific website usage data is aggregated using trusted execution environment computing hardware. This ensures that privacy is preserved while user-specific data is transferred to and processed by the system. The trusted execution environment applies differential privacy techniques to ensure that use of the aggregated information does not reveal individually identifying information about a user's website usage history. In this way, privacy is maintained while still enabling many of the scenarios that would otherwise rely on third-party cookies.

(7) Differential privacy allows data to be shared and analyzed in a way that provides strong privacy protections for individuals while still allowing meaningful insights to be drawn from website usage data. Differential privacy is a statistical method that obfuscates the output of a function that operates on user-specific data. The output is obfuscated such that it is impractical to determine if a given user record was present in the output data. Specifically, given two datasets, one including a specific user's data and another without, differential privacy ensures that with high probability it cannot be determined which dataset includes the specific user's data. For example, user-specific data that tracks views of cat videos is differentially private if it can't be determined whether a user's cat video viewing history is part of one of two different datasets. Useful information may still be derived from a differentially private function, such as the average number of cat videos viewed. While differential privacy is the most commonly discussed standard of privacy discussed herein,

other privacy standards are similarly contemplated.

(8) Differential privacy is transitive—a function applied to the output of a differentially private function is also differentially private. In this way, once differential privacy has been introduced into a sequence of processing steps, the results remain differentially private throughout the sequence. In order to protect user privacy of website usage information before it has been processed into a differentially private form, a trusted execution environment is employed.

(9) A trusted execution environment uses encryption to ensure that website usage data is secure while it is aggregated and processed into a differentially private dataset. For example, a trusted execution environment may utilize a public/private key pair to ensure that user-specific website usage data is received securely from a web browser. In some configurations, the trusted execution environment is implemented by processors that are each assigned a private key. Website usage data may be encrypted with the corresponding public key before being transferred to the trusted execution environment. In this way, only code executing within the trusted execution environment is capable of accessing the user-specific website usage data.

(10) The trusted execution environment may also use the processor-specific public/private key pair to attest to a third-party attestation service that it is running a specific piece of code and/or that the execution environment is in fact trusted. In this scenario, once the third-party attestation service has verified the trusted execution environment, the third-party attestation service provides a public/private key pair to the trusted execution environment for use encrypting website usage data. For example, the trusted execution environment may provide the public key obtained from the third-party attestation service to a client device so that the client device is able to encrypt website usage data before transmitting it to the trusted execution environment. Other TEEs may verify with the third-party attestation service that they are in fact trusted, e.g., that they are running secure software on secure hardware. These TEEs may be given the same public/private key pair, enabling them to further process data that was received previously by a different TEE.

(11) Features and technical benefits other than those explicitly described above will be apparent from a reading of the following Detailed Description and a review of the associated drawings. This Summary is provided to introduce a selection of concepts in a simplified form that are further described below in the Detailed Description. This Summary is not intended to identify key or essential features of the claimed subject matter, nor is it intended to be used as an aid in determining the scope of the claimed subject matter. The term “techniques,” for instance, may refer to system(s), method(s), computer-readable instructions, module(s), algorithms, hardware logic, and/or operation(s) as permitted by the context described above and throughout the document.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The Detailed Description is described with reference to the accompanying figures. In the figures, the left-most digit(s) of a reference number identifies the figure in which the reference number first appears. The same reference numbers in different figures indicate similar or identical items. References made to individual items of a plurality of items can use a reference number with a letter of a sequence of letters to refer to each individual item. Generic references to the items may use the specific reference number without the sequence of letters.

(2) FIG. 1 illustrates using third-party cookies to serve targeted content to a user.

(3) FIG. 2 illustrates an overview of a system for privacy preserving user personalization using noisy ranking.

(4) FIG. 3 illustrates a system for privacy preserving user personalization using noisy ranking.

(5) FIG. 4 is a flow diagram of an example method for privacy preserving user personalization using noisy ranking.

(6) FIG. 5 is a computer architecture diagram illustrating an illustrative computer hardware and software architecture for a computing system capable of implementing aspects of the techniques and technologies presented herein.

(7) FIG. 6 is a diagram illustrating a distributed computing environment capable of implementing aspects of the techniques and technologies presented herein.

DETAILED DESCRIPTION

(8) Cookies are small text files that a website can place on a user's computer or mobile device. When the user visits the website again, the website can retrieve the cookie and use the information therein to remember the user's preferences, login information, browsing history, etc. For example, a website may use a cookie to store how long a user visited a particular page, what page the user visited next, etc. In the context of an advertising platform, the cookie may store indications of which ads were viewed, clicked on, or otherwise interacted with. In some configurations, the website may store a user identifier in the cookie while other user-specific information is stored on the server. Then on subsequent visits the website may retrieve the identifier from the cookie and use it to retrieve the user-specific information from the server.

(9) As discussed above, third-party cookies raise privacy concerns. One concern is that by tracking user activity across a number of websites, a single entity is able to obtain a significant amount of information about an individual user. Information collected by an entity in this way is only as private as the security of the entity's systems and the willingness of the entity to protect the user's privacy.

(10) For example, an advertising platform may use third-party cookies to track website usage patterns across websites. This allows sensitive information to be obtained about identifiable users. A breach of the advertising service's systems may result in user data being released to the public. Malicious entities could leverage this data to solicit users or even blackmail them.

(11) FIG. 1 illustrates how third-party cookies may be used to serve targeted content to a user. Targeted content refers to the practice of tailoring and delivering specific information, advertisements, or recommendations to users based on their personal data, preferences, or online behavior. The aim is to provide users with content that is more relevant, engaging, and likely to generate a positive response, such as clicking on an ad or purchasing a product. Serving targeted content, such as news articles, videos, or advertisements, etc., is one example of a technique that leverages information gathered with third-party cookies.

(12) As illustrated, user **102** operates computing device **104**, navigating browser **106** to publisher website **110**. Publisher website **110** has domain name **112**, which is part of a uniform resource locator (URL). Browser **106** may implement a same-origin policy, a security feature that restricts how documents and scripts from one domain can interact with resources from another domain. Unless special permission is granted by the user, this constrains website **110** to interacting with cookies that originated from domain name **112**. User **102** has provided username **113**, such as when logging into website **110**.

(13) Website **110** includes third-party component **120**. Third-party component **120** may be an embedded website, a web component, a web application, a JavaScript library, or any other type of web content provided by content provider **108**. Third-party component **120** has domain name **114** associated with content provider **108**. When third-party cookies are enabled in browser **106**, third-party component **120** is enabled to store data in and read data from third-party cookie **122** because they are both associated with the same domain name **114**. Third-party cookie **122** may also be used to store the user-specific data that has been gathered from browser **106**. For example, clicks **126**, views **128**, and time spent **129** are examples of the types of data that may be stored in third-party cookies, and are described in additional detail below. Such user-specific data may be used by content provider **108** to refine and provide targeted content.

(14) Views **128** are indications of views of a website or a portion of a website such as a news article or an ad. In some configurations, a view is recorded if a website or a portion of a website is

displayed by computing device **104**. Views may also be identified based on eye tracking, monitoring mouse cursor position, or other more active techniques for determining if a piece of content has been viewed. Time **129** indicates how long a web page was viewed. Time **129** may refer to an amount of time that user **102** was actively using website **110**, the amount of time that website was inactive but still displayed and/or open, or a combination thereof. Clicks **126** record user interactions with website **110** and web pages therein. For example, clicks **126** may include a history of mouse movements, mouse button clicks, mouse over events, mouse button click and holds, and the like. Logged clicks **166** represents the submission of clicks **126** to impression and click store **138**.

(15) In this example, key **124**, which uniquely identifies user **102**, is also stored in third-party cookie **122**. Key **124** may have been provided by content provider **108** to track website usage of user **102** across multiple websites. Key **124** may be any type of unique identifier, such as a string of numbers, a username, a globally unique identifier, etc. Key **124** may be used when submitting user-specific data that has been gathered from browser **106**. For example, async update **160** provides data stored in third-party cookie **122** to individual user data store **130**. Key **124** may also be used when requesting content, e.g., as part of request **162**.

(16) Content provider **108** leverages a variety of data stores and machine learning models to process content requests **162**. As illustrated, content provider **108** contains or is associated with individual user data store **130**, aggregate user data tables **132**, user models **134**, content provider data **136**, and impression and click store **138**. Impression and click store **138** stores interactions with specific pieces of content that have been incorporated into a publisher's website.

(17) In order to serve a piece of content in response to content request **162**, content provider **108** may perform the steps of creating a user representation **140**, retrieving relevant pieces of content **142**, and ranking and selecting content **144**. One or more pieces of selected content **164** are then returned to publisher website **110**.

(18) Individual user data store **130** is a cloud-hosted store of individualized user data obtained from third-party component **120**, and which may have been stored for at least a period of time in third-party cookie **122**. Individual user data store **130** may be associated with key **124**, and as such contains potentially sensitive website usage information that can be tied to a particular user. Individual data store **130** tracks the websites that a user visits, as a way of building an understanding of user browsing habits.

(19) Aggregate user data tables **132** represent user data that has been processed by content provider **108**. The create user data tables operation **150** creates and/or updates aggregate user data tables **132** based in part on website usage data stored in individual user data store **130** and impression and click store **138**. Aggregate user data tables **132** include information used to perform the create user representation operation **140**.

(20) User models **134** are machine learning models trained on website usage information. User models **134** are created by the train models operation **152** based in part on impressions and clicks stored in impression and click store **138**. These models may be updated periodically, from every few minutes to every day or even longer periods of time. User models **134** may also be used to perform the create user representation operation **140**.

(21) Profile **141** is created by the create user representation operation **140**. Profile **141** is a record of user **102** used to select targeted content, such as a targeted advertisement. User profile **141** may contain demographic information, website usage information, and/or attributes derived therefrom.

(22) The retrieve relevant pieces of content operation **142** uses profile **141** to retrieve one or more pieces of content that are best suited to be presented to user **102**. Content may be retrieved from content provider data **136**. Content may be selected based on a number of factors, including a predicted likelihood of user **102** interacting with a piece of content, a profit or other value associated with displaying a particular piece of content, or the like. In the context of an advertising platform, a predicted conversion rate is one factor that may be used in selecting relevant pieces of

content **143**. Rank and select content operation **144** selects one or more of relevant pieces of content **143** to be returned to website **110** as selected content **164**. In some configurations, rank and select content operation **144** performs an auction in which selected content **164** is selected from multiple competitive suppliers of content. This selection may be made according to budget, targeting, and campaign objectives of content providers, as well as the prices set by publishers.

(23) FIG. 2 illustrates an overview of a system for privacy preserving content delivery using noisy ranking. FIG. 2 is a conceptual overview—more details are provided below in conjunction with FIG. 3. User **102**, client device **104**, browser **106**, publisher website **110**, and domain name **112** are as describe above in conjunction with FIG. 1, except that in FIG. 2 third-party cookies have been disabled by browser **106**. Third-party component **220** is a component similar to component **120**. For example, third-party component **220** may be a web component published by an advertising platform. Third-party component **220** may be embedded in publisher website **110** where it observes user behavior and displays advertisements.

(24) Third-party component **220** sends user-specific website usage data **260** to serving engine **202**, instead of tracking user **102** with third-party cookies **122** as in FIG. 1. This user-specific website usage data **260** may include sensitive, private information of user **102**. Therefore, in the illustrated example, the user-specific website usage data **260** is protected by protected environment **200** until it can be processed into privacy-enhanced data tables, models, and other privacy-preserving forms.

(25) Encryption is used to protect usage data **260** as it is transmitted from browser **106** to protected environment **200**. Serving engine **202** contains hardware that ensures the privacy of web usage data **260** within protected environment **200**. Specifically, serving engine **202** executes on trusted execution environment hardware. For example, serving engine **202** may execute on special hardware with a private key that makes it uniquely able to decrypt the encrypted usage data **260**. The trusted execution environment may also include a hardware-based feature that attests that the code it is executing has been vetted by an independent third party. For example, the trusted execution environment may attest that it is running code that has been published as open source and that has been vetted by an independent security analyst that usage data **260** will not be released outside of the protected environment **200**. The trusted execution environment may also use public/private keys to attest that the TEE is using agreed upon hardware and software.

(26) In some configurations, instead of using the TEE's public/private key pair to encrypt usage data **260**, a trusted third-party may use a different public/private key pair to encrypt usage data **260**. Then, only after TEE has established that it running appropriately secure software and hardware will the trusted third-party provide the TEE with the key to decode usage data **260**.

(27) Aggregation engine **204** and/or model training engine **206** processes the website usage data **260** into low-dimensional embeddings. A low-dimensional embedding is a mathematical representation of high-dimensional data in a lower-dimensional space. It is a technique used in machine learning, data analysis, and visualization to simplify complex datasets and reveal patterns, relationships, or structures that may be difficult to discern in the original high-dimensional space. Specifically, high-dimensional website usage data **260** is converted into a lower-dimensional space. An ad-user relevance ranking can then be obtained by performing an inner product between the low-dimensional user embedding and ad embeddings. Ad embeddings with larger dot products are considered more relevant than ad embeddings with small dot products.

(28) Low-dimensional user data may also be provided to aggregation engine **204**. In order to protect the anonymity of the user, aggregation engine **204** may employ privacy-enhanced techniques such as differential privacy on the data it generates. For example, aggregation engine **204** may aggregate usage data **260** for user **102** with usage data from other users. This reduces the likelihood that someone could determine that usage data **260** from the particular user **102** was included in the aggregation. Noise—random changes to the data—may also be incorporated to prevent user **102** from being identified as part of the dataset. For example, values may be adjusted up or down at random at a rate that obscures whether a particular user's data was included in the

aggregation. Noise may also be added by replacing values from usage data **260** with randomly generated values. In differential privacy, noise is typically generated from probability distributions like the Laplace distribution or the Gaussian distribution. The amount of noise added is determined by the sensitivity of the function being computed—the maximum difference in the output when a single record is added or removed from the dataset. Model training engine **206** may also train machine learning models that observe privacy-enhancing techniques such as differential privacy. (29) Differentially private data yielded by aggregation engine **204** and model training engine **206** may be reinserted **242** to serving engine **202** in order to serve targeted content while preserving user privacy. The differentially private data yielded by aggregation engine **204** and model training engine **206** may also be released **240** for use by components outside of protected environment **200**. For example, a component outside the protected environment **200** that computes ad prices may utilize differentially private aggregations of website usage data and/or differentially private machine learning models trained on website usage data.

(30) Serving engine **202** may also receive a request **262** for content, such as a request for an advertisement. Serving engine **202** may use differentially private data tables generated by aggregation engine **204** and/or differentially private machine learning models trained by model training engine **206** to select content. Serving engine **202** may use usage data **260** received from browser **106** with differentially private data tables and differentially private models to make this selection, e.g., by taking dot products of an embedding generated from usage data **260** and embeddings of the differentially private tables/models. Embeddings of the differentially private tables/models associated with the highest dot products are considered to be a better match for the usage data **260**. Serving engine **202** may also consider additional data **230** when selecting content. Additional data **230** is not derived from third-party website usage data.

(31) FIG. 3 illustrates a system for privacy preserving content delivery using noisy ranking. FIG. 3 illustrates one example of an implementation of some of the concepts depicted in FIG. 2. User **102**, client device **104**, browser **106**, and publisher website **110** are as described above in conjunction with FIGS. 1 and 2. In some configurations, individual user data store **330** stores user data locally, instead of on individual user data store **130** as in FIG. 1. This increases privacy, as user data is no longer vulnerable to leaks or hacks of individual user data store **130**. Data stored in individual user data store **330** is encrypted and transmitted to a trusted execution environment within the protected environment **200** as part of content request **362**. Furthermore, content request **362** and logged clicks **366** are encrypted using public key **325**, which is selected based on the specific computing device that will receive the content request and logged clicks. In some configurations, content request and logged clicks are encrypted using public key **392**, which may be obtained from trusted execution environment **341**. TEE **341** may have obtained public key **392** from attestation service **390** after using private key **360** to attest that TEE **341** is in fact a trusted execution environment.

(32) Create user representation operation **340**, retrieve relevant pieces of content operation **342**, rank and select content operation **344**, create user data tables operation **350**, and train models operation **352** may be implemented by one or more trusted execution environments (TEEs) such as trusted execution environment **341**. A TEE is a secure area of a computer or mobile device that ensures sensitive data is processed securely. It is typically isolated from the main operating system and runs on a separate processor or core. The TEE provides a secure environment for cryptographic key storage, ensuring that keys are protected from malware and other forms of tampering.

(33) Some trusted execution environments include processors manufactured with a private key that is stored internally, without any means for copying the key. In these embodiments, each processor is associated with a unique private key. The processor's manufacturer may delete all other copies of these private keys before releasing the processor to market, ensuring that data encrypted with the corresponding public keys may only be accessed by the appropriate processor.

(34) The manufacturer of the trusted execution environment may also publish table **322** or other data source that lists processor identifiers and their corresponding public keys. Third-party

component **220** may receive a processor ID **326** when connecting to content provider **108**. For example, third-party component **220** may receive an indication that TEE **341** has processor ID **1**. Third-party component **220** may also receive processor ID **326** when submitting usage data **260**, when submitting request **262** or at any other time. Third party component **220** may use this processor ID to obtain public key **323** from server-key table **322**. In some configurations, before sending request **362**, logged clicks **366**, or any other request to content provider **108**, third-party component **220** may use this public key **323** to encrypt some or all of the request.

(35) In other configurations, an attestation service **390** may independently verify that TEE **341** is running code that has been deemed secure by an independent party and/or that TEE **341** uses hardware to effectively protect usage data **260**. Once attestation service **390** has established the veracity of TEE **341**, attestation service **390** may transmit public key **392** and private key **394** to key administrator **398**. Key administrator **398** may be within TEE **341** or separate from TEE **341**, and key administrator **398** may be within or external to protected environment **200**.

(36) Key administrator **398** may provide third-party component **220** with public key **392**. Third-party component **220** may use public key **392** to encrypt usage data **260** before transmitting it to TEE **341**. TEE **341** may receive private key **394** from attestation service key administrator **398**. TEE **341** may then use private key **394** to access usage data **260** that was encrypted with public key **392**, enabling usage data **260** to be processed in secure environment **200**.

(37) TEE **341** may transmit a report to another TEE within secure environment **200**. If the report is not differentially private then privacy may be maintained by encrypting the report with public key **392**. The other TEE will need to repeat attestation with attestation service **390** in order to acquire private key **394** with which to decrypt the report. This technique may be used again with another TEE, enabling a sequence of TEEs to operate on usage data **260**.

(38) In some configurations, multiple attestation services **390** may be employed to provide multiple public/private key pairs. Usage data **260** may then be doubly encrypted, ensuring that privacy is maintained even if one of the key pairs is compromised. Additionally, in some configurations the Key Administrators **398** are separate entities from the attestation services **390**, simply providing private keys to TEEs that attestation service **390** attests are properly configured and secure. Such configurations introduce more required points of failure before user privacy can be compromised.

(39) These techniques and equivalent techniques protect potentially sensitive web usage data while it is transmitted from browser **106** to trusted execution environment **341** within the protected environment **200** of content provider **108**.

(40) Once request **362** is received by trusted execution environment **341**, the create user representation operation **340** is performed, followed by the retrieve relevant pieces of content operation **342** and the rank and select content operation **344**, after which selected content **364** is returned to third-party component **220** for display in browser **106**. While create user representation operation **340**, retrieve relevant pieces of content operation **342**, and rank and select content operation **344** are illustrated in FIG. 3, these are simply examples of operations that may be applied to return selected content **364** in response to request **362**. Other operations are similarly contemplated. Furthermore, while FIG. 3 illustrates one application of the disclosed privacy preserving user tracking techniques—selecting content—other applications such as web analytics are similarly contemplated.

(41) Request **362** is decrypted by trusted execution environment **341** using the private key **360** that corresponds to the public key **323** listed for processor ID '1' in server-key table **322**. Additionally, or alternatively, request **362** may be decrypted by TEE **341** using private key **394** obtained from the key administrator **398**. Once decrypted, request **362** may be processed in order to return one or more appropriate pieces of content to third-party component **220**. Request **362** may include, for example, web usage data from individual user data **330**. This web usage data is potentially detailed and sensitive, but is confined to trusted execution environment **341**.

(42) As illustrated, each of processing steps **340**, **342**, and **344** are executed in a different trusted

execution environment. In this scenario, communications between each trusted execution environment may be encrypted similar to how request **362** is encrypted. For example, the public key of the destination trusted execution environment may be obtained from server-key table **322** and used to encrypt the communication. For example, user representation **380**, which may be derived from sensitive web usage data, may be encrypted when passing it from the create user representation operation **340** to the retrieve relevant pieces of content operation **342**. In this way, even if they are obtained, communications within protected environment **200** remain unusable to unauthorized users. In another configuration, each TEE used to perform one of processing steps **340-344** may obtain public key **392** and private key **394**, which originate from service **390**, and each TEE may use these keys to encrypt and decrypt data sent between TEEs.

(43) In some configurations, when training models **352** and/or aggregating user data tables **350**, a differential privacy administrator ensures that differential privacy is maintained by enforcing constraints on how differentially private models **352** and/or data tables **350** are generated. The differential privacy administrator may be implemented by another trusted execution environment or a trusted third-party service. The differential privacy administrator receives or generates metadata about records used to train the models **352** and/or aggregate tables **350**, such as a hash of a record. One example of a constraint uses the hash of the record to ensure that the record is only processed once, or some other defined number of times. This maintains differential privacy because differential privacy is not effective if an unrestricted amount of noise is added to models **352** and/or tables **350**.

(44) Create user representation operation **340**, retrieve relevant pieces of content operation **342**, and rank and select content operation **344** operate similar to operations **140**, **142**, and **144** described above in conjunction with FIG. **1**, with the addition in FIG. **3** of the use of encryption to secure the privacy of user **102**. One difference is that user data is no longer stored on a server device such as individual user data store **130** in FIG. **1**, but is instead stored in individual user data **330** within browser **106**. Another difference is that the information provided by aggregate user data tables **332** and user models **334** is differentially private. This means that potentially sensitive web usage data has been modified by the addition of noise to make it impracticable to determine whether a profile of a user was part of a particular data table or used to train a particular model.

(45) Create user representation operation **340** may use individual user data **330** received from request **362** to create a representation **380** of user **102**. This representation **380** may indicate, for example, that user **102** has viewed a large number of cat videos.

(46) Retrieve relevant pieces of content operation **342** may utilize user representation **380** to obtain content from content provider data store **136**. For example, retrieve relevant pieces of content operation **342** may obtain a list of videos, advertisements, or other content based on user representation **380**. In some configurations, content candidates **382** are obtained by taking an inner product of an embedding of user representation **380** and embeddings of aggregate user data tables **332** and/or one or more of user models **334**. Content may be selected to be provided to third-party component **220** based on a ranking of the inner products, with content that has a larger inner product being ranked higher than content with a lower inner product. In some configurations, retrieve relevant pieces of content operation **342** may perform a real-time differentially private request of aggregate user data tables **332**, user models **334**, or any other differentially private data store.

(47) If rank and select content operation **344** operates in a different trusted execution environment, then content candidates **382** may be encrypted using the public key of the destination TEE. When a third-party attestation service **390** is used, then content candidates **382** may be encrypted with the public key **392** obtained from third-party attestation service **390**. Rank and select content operation **344** then ranks, filters, and otherwise evaluates the content candidates **382**. One or more of the top-ranked pieces of content are then obtained from content store **339**. One or more of these pieces of content are then returned to third-party component **220** as selected content **364**.

(48) In some configurations, a k-anonymity check is applied to ensure that selected content **364** does not reveal a user's personal information or online behavior. K-anonymity is a property of datasets that aims to protect the privacy of individuals within the dataset by ensuring that each individual is indistinguishable from at least $k-1$ other individuals. This is achieved by generalizing or suppressing certain attributes (also known as quasi-identifiers) in the dataset to create groups of records that are indistinguishable from each other. By ensuring that selected content **364** adheres to k-anonymity standards, individual user data **330** may be used to select content while ensuring that individual users cannot be singled out and their privacy is maintained.

(49) As illustrated, rank and select content operation **344** also produces impressions **302**. While selected content **364** is returned to browser **106** for display, impressions **302** are used to further train user models **334** and update aggregate user data tables **332**. Impressions **302** may include reference to or copies of selected content **364**. In the case of an advertising platform, an impression indicates that user **102** has been presented with an advertisement. In some configurations, selected content **364** is considered an impression. Impressions **302** and logged clicks **366** are stored in impression and click store **338**. In some configurations, impressions **302** and logged clicks **366** are encrypted when stored in impression and click store **338**.

(50) Create user data tables operation **350** queries impression and click store **338** to generate an aggregate user data table that is stored in aggregate user data tables **332**. Create user data tables operation **350** generates tables that are differentially private. For example, noise—random fluctuations in values—may be introduced to obscure whether web usage data from user **102** is or is not in one of the generated tables. Typically, for tables with larger numbers of elements, less noise needs to be introduced to maintain differential privacy. Similarly, train models operation **352** uses data from impression and click store **338** to create a differentially private machine learning model stored in user models **334**. These models may also be used to select content targeted to a particular user.

(51) In FIG. 3, the larger dashed lines between stores and components indicate that data is differentially private. Solid lines indicate that the communication is encrypted. Shorter dashed lines indicate the boundary of protected environment **200**. Since an operation performed on the output of a differentially private function is also differentially private, and create user data tables **350** and train models **352** introduce differentially private data into the system, downstream of these operations, the data the operations produced may be used outside of the protected environment without eroding user privacy.

(52) The system depicted in FIG. 3 illustrates a loop—web usage data is securely gathered from browsers and used to generate data tables and models. These tables and models are used to select content. The content that is selected is an impression, which is stored in impression and click store **338**, completing the loop. By ensuring that user data tables and models are differentially private, the incorporation of this data into subsequent iterations of the loop ensures user privacy.

(53) With reference to FIG. 4, routine **400** begins at operation **402**, where encrypted, user-specific web usage data **366** is received by a trusted execution environment. In some configurations, the user-specific web usage data **366** has been encrypted with a public key **392** obtained from a third-party attestation service **390**.

(54) Next at operation **404**, the web-usage data **366** is decrypted within protected environment **200** using private key **362** stored in trusted execution environment **343**.

(55) Next at operation **406**, a differentially private machine learning model **334** is trained based on the user-specific web usage data **366**.

(56) Next at operation **408**, a content request **362** is received from a third-party component **320** running on browser **106** of client device **104**.

(57) Next at operation **410**, content **364** is selected based on the differentially private machine learning model **334**.

(58) Then, at operation **412**, the selected content **364** is checked for k-anonymity before being

returned to the third-party component **320**.

(59) The particular implementation of the technologies disclosed herein is a matter of choice dependent on the performance and other requirements of a computing device. Accordingly, the logical operations described herein are referred to variously as states, operations, structural devices, acts, or modules. These states, operations, structural devices, acts, and modules can be implemented in hardware, software, firmware, in special-purpose digital logic, and any combination thereof. It should be appreciated that more or fewer operations can be performed than shown in the figures and described herein. These operations can also be performed in a different order than those described herein.

(60) It also should be understood that the illustrated methods can end at any time and need not be performed in their entirety. Some or all operations of the methods, and/or substantially equivalent operations, can be performed by execution of computer-readable instructions included on a computer-storage media, as defined below. The term “computer-readable instructions,” and variants thereof, as used in the description and claims, is used expansively herein to include routines, applications, application modules, program modules, programs, components, data structures, algorithms, and the like. Computer-readable instructions can be implemented on various system configurations, including single-processor or multiprocessor systems, minicomputers, mainframe computers, personal computers, hand-held computing devices, microprocessor-based, programmable consumer electronics, combinations thereof, and the like.

(61) Thus, it should be appreciated that the logical operations described herein are implemented (1) as a sequence of computer implemented acts or program modules running on a computing system and/or (2) as interconnected machine logic circuits or circuit modules within the computing system. The implementation is a matter of choice dependent on the performance and other requirements of the computing system. Accordingly, the logical operations described herein are referred to variously as states, operations, structural devices, acts, or modules. These operations, structural devices, acts, and modules may be implemented in software, in firmware, in special purpose digital logic, and any combination thereof.

(62) For example, the operations of the routine **400** are described herein as being implemented, at least in part, by modules running the features disclosed herein can be a dynamically linked library (DLL), a statically linked library, functionality produced by an application programming interface (API), a compiled program, an interpreted program, a script or any other executable set of instructions. Data can be stored in a data structure in one or more memory components. Data can be retrieved from the data structure by addressing links or references to the data structure.

(63) Although the following illustration refers to the components of the figures, it should be appreciated that the operations of the routine **400** may be also implemented in many other ways. For example, the routine **400** may be implemented, at least in part, by a processor of another remote computer or a local circuit. In addition, one or more of the operations of the routine **400** may alternatively or additionally be implemented, at least in part, by a chipset working alone or in conjunction with other software modules. In the example described below, one or more modules of a computing system can receive and/or process the data disclosed herein. Any service, circuit or application suitable for providing the techniques disclosed herein can be used in operations described herein.

(64) FIG. 5 shows additional details of an example computer architecture **500** for a device, such as a computer or a server configured as part of the systems described herein, capable of executing computer instructions (e.g., a module or a program component described herein). The computer architecture **500** illustrated in FIG. 5 includes processing unit(s) **502**, a system memory **504**, including a random-access memory **506** (“RAM”) and a read-only memory (“ROM”) **508**, and a system bus **510** that couples the memory **504** to the processing unit(s) **502**.

(65) Processing unit(s), such as processing unit(s) **502**, can represent, for example, a CPU-type processing unit, a GPU-type processing unit, a field-programmable gate array (FPGA), another

class of digital signal processor (DSP), or other hardware logic components that may, in some instances, be driven by a CPU. For example, and without limitation, illustrative types of hardware logic components that can be used include Application-Specific Integrated Circuits (ASICs), Application-Specific Standard Products (ASSPs), System-on-a-Chip Systems (SOCs), Complex Programmable Logic Devices (CPLDs), etc.

(66) A basic input/output system containing the basic routines that help to transfer information between elements within the computer architecture **500**, such as during startup, is stored in the ROM **508**. The computer architecture **500** further includes a mass storage device **512** for storing an operating system **514**, application(s) **516**, modules **518**, and other data described herein.

(67) The mass storage device **512** is connected to processing unit(s) **502** through a mass storage controller connected to the bus **510**. The mass storage device **512** and its associated computer-readable media provide non-volatile storage for the computer architecture **500**. Although the description of computer-readable media contained herein refers to a mass storage device, it should be appreciated by those skilled in the art that computer-readable media can be any available computer-readable storage media or communication media that can be accessed by the computer architecture **500**.

(68) Computer-readable media can include computer-readable storage media and/or communication media. Computer-readable storage media can include one or more of volatile memory, nonvolatile memory, and/or other persistent and/or auxiliary computer storage media, removable and non-removable computer storage media implemented in any method or technology for storage of information such as computer-readable instructions, data structures, program modules, or other data. Thus, computer storage media includes tangible and/or physical forms of media included in a device and/or hardware component that is part of a device or external to a device, including but not limited to random access memory (RAM), static random-access memory (SRAM), dynamic random-access memory (DRAM), phase change memory (PCM), read-only memory (ROM), erasable programmable read-only memory (EPROM), electrically erasable programmable read-only memory (EEPROM), flash memory, compact disc read-only memory (CD-ROM), digital versatile disks (DVDs), optical cards or other optical storage media, magnetic cassettes, magnetic tape, magnetic disk storage, magnetic cards or other magnetic storage devices or media, solid-state memory devices, storage arrays, network attached storage, storage area networks, hosted computer storage or any other storage memory, storage device, and/or storage medium that can be used to store and maintain information for access by a computing device.

(69) In contrast to computer-readable storage media, communication media can embody computer-readable instructions, data structures, program modules, or other data in a modulated data signal, such as a carrier wave, or other transmission mechanism. As defined herein, computer storage media does not include communication media. That is, computer-readable storage media does not include communications media consisting solely of a modulated data signal, a carrier wave, or a propagated signal, per se.

(70) According to various configurations, the computer architecture **500** may operate in a networked environment using logical connections to remote computers through the network **520**. The computer architecture **500** may connect to the network **520** through a network interface unit **522** connected to the bus **510**. The computer architecture **500** also may include an input/output controller **524** for receiving and processing input from a number of other devices, including a keyboard, mouse, touch, or electronic stylus or pen. Similarly, the input/output controller **524** may provide output to a display screen, a printer, or other type of output device.

(71) It should be appreciated that the software components described herein may, when loaded into the processing unit(s) **502** and executed, transform the processing unit(s) **502** and the overall computer architecture **500** from a general-purpose computing system into a special-purpose computing system customized to facilitate the functionality presented herein. The processing unit(s) **502** may be constructed from any number of transistors or other discrete circuit elements,

which may individually or collectively assume any number of states. More specifically, the processing unit(s) **502** may operate as a finite-state machine, in response to executable instructions contained within the software modules disclosed herein. These computer-executable instructions may transform the processing unit(s) **502** by specifying how the processing unit(s) **502** transition between states, thereby transforming the transistors or other discrete hardware elements constituting the processing unit(s) **502**.

(72) FIG. **6** depicts an illustrative distributed computing environment **600** capable of executing the software components described herein. Thus, the distributed computing environment **600** illustrated in FIG. **6** can be utilized to execute any aspects of the software components presented herein. For example, the distributed computing environment **600** can be utilized to execute aspects of the software components described herein.

(73) Accordingly, the distributed computing environment **600** can include a computing environment **602** operating on, in communication with, or as part of the network **604**. The network **604** can include various access networks. One or more client devices **606A-606N** (hereinafter referred to collectively and/or generically as “clients **606**” and also referred to herein as computing devices **606**) can communicate with the computing environment **602** via the network **604**. In one illustrated configuration, the clients **606** include a computing device **606A** such as a laptop computer, a desktop computer, or other computing device; a slate or tablet computing device (“tablet computing device”) **606B**; a mobile computing device **606C** such as a mobile telephone, a smart phone, or other mobile computing device; a server computer **606D**; and/or other devices **606N**. It should be understood that any number of clients **606** can communicate with the computing environment **602**.

(74) In various examples, the computing environment **602** includes servers **608**, data storage **610**, and one or more network interfaces **612**. The servers **608** can host various services, virtual machines, portals, and/or other resources. In the illustrated configuration, the servers **608** host virtual machines **614**, Web portals **616**, mailbox services **618**, storage services **620**, and/or, social networking services **622**. As shown in FIG. **6** the servers **608** also can host other services, applications, portals, and/or other resources (“other resources”) **624**.

(75) As mentioned above, the computing environment **602** can include the data storage **610**. According to various implementations, the functionality of the data storage **610** is provided by one or more databases operating on, or in communication with, the network **604**. The functionality of the data storage **610** also can be provided by one or more servers configured to host data for the computing environment **602**. The data storage **610** can include, host, or provide one or more real or virtual datastores **626A-626N** (hereinafter referred to collectively and/or generically as “datastores **626**”). The datastores **626** are configured to host data used or created by the servers **608** and/or other data. That is, the datastores **626** also can host or store web page documents, word documents, presentation documents, data structures, algorithms for execution by a recommendation engine, and/or other data utilized by any application program. Aspects of the datastores **626** may be associated with a service for storing files.

(76) The computing environment **602** can communicate with, or be accessed by, the network interfaces **612**. The network interfaces **612** can include various types of network hardware and software for supporting communications between two or more computing devices including, but not limited to, the computing devices and the servers. It should be appreciated that the network interfaces **612** also may be utilized to connect to other types of networks and/or computer systems.

(77) It should be understood that the distributed computing environment **600** described herein can provide any aspects of the software elements described herein with any number of virtual computing resources and/or other distributed computing functionality that can be configured to execute any aspects of the software components disclosed herein. According to various implementations of the concepts and technologies disclosed herein, the distributed computing environment **600** provides the software functionality described herein as a service to the computing

devices. It should be understood that the computing devices can include real or virtual machines including, but not limited to, server computers, web servers, personal computers, mobile computing devices, smart phones, and/or other devices. As such, various configurations of the concepts and technologies disclosed herein enable any device configured to access the distributed computing environment **600** to utilize the functionality described herein for providing the techniques disclosed herein, among other aspects.

(78) The present disclosure is supplemented by the following example clauses:

(79) Example 1: A method comprising: establishing, to a third-party attestation service (**390**), that the processing system (**502**) is a trusted execution environment; receiving a public key (**392**) and a corresponding private key (**394**) from the third-party attestation service (**390**); receiving user-specific web usage data (**366**) that has been encrypted with the public key (**392**); decrypting the user-specific web usage data (**366**) using the private key (**394**); storing the user-specific web usage data (**366**) in a differentially private data store (**332**); selecting content (**364**) to be displayed to a user (**102**) based on the differentially private data store (**332**); and transmitting the selected content (**364**) to a client computing device (**104**).

(80) Example 2: The method of Example 1, wherein the differentially private data store comprises a differentially private machine learning model or a differentially private data table.

(81) Example 3: The method of Example 2, further comprising: providing the differentially private machine learning model and the differentially private data table to an external service.

(82) Example 4: The method of Example 2, further comprising: providing the differentially private machine learning model and the differentially private data table to an internal service that selects content for users; and deleting the non-differentially private user-specific web usage data.

(83) Example 5: The method of Example 1, further comprising: performing a k-anonymity check on the content to be displayed to remove any identifying information within the content or within a reference to the content.

(84) Example 6: The method of Example 1, further comprising: providing the client computing device with the public key received from the third-party attestation service.

(85) Example 7: The method of Example 2, wherein a publicly disclosed program executing on the trusted execution environment trains the differentially private machine learning model.

(86) Example 8: A computer-readable storage medium (**504**) having computer-executable instructions stored thereupon that, when executed by a processing system (**502**), cause the processing system (**502**) to: establish, to a third-party attestation service (**390**), that the processing system (**502**) is a trusted execution environment; receive a public key (**392**) and a corresponding private key (**394**) from the third-party attestation service (**390**); receive user-specific web usage data (**366**) generated by a user (**102**) and that has been encrypted with the public key (**392**); decrypt the user-specific web usage data (**366**) using the private key (**394**) of the trusted third-party attestation service (**390**); within the trusted execution environment (**343**), train a differentially private machine learning model (**334**) with the user-specific web usage data (**366**); receive a content request (**362**) from a client computing device (**104**); select content (**364**) to be displayed to the user (**102**) based on the differentially private machine learning model (**334**); encrypt the selected content (**364**); and transmit the encrypted selected content (**364**) to the client computing device (**104**) in response to the content request (**362**).

(87) Example 9: The computer-readable storage medium of Example 8, wherein individual user data is stored on the client computing device before being processed by the trusted execution environment.

(88) Example 10: The computer-readable storage medium of Example 8, wherein the content is selected from a content store that encrypts the content.

(89) Example 11: The computer-readable storage medium of Example 8, wherein the trusted-execution environment establishes that the processing system is a trusted execution environment based on a private key that is permanently stored in the processor of the processing system.

(90) Example 12: The computer-readable storage medium of Example 8, further comprising: establishing an attested channel that is used to communicate the user-specific web usage data to the trusted execution environment.

(91) Example 13: The computer-readable storage medium of Example 8, wherein the trusted execution environment is one of a plurality of trusted execution environments, each having a unique public key/private key pair received from the trusted third-party attestation service, and wherein the plurality of trusted execution environments perform a plurality of operations that collectively select the content to be displayed to the user.

(92) Example 14: The computer-readable storage medium of Example 13, wherein the plurality of trusted execution environments comprise a sequence of trusted execution environments, and wherein each of the sequence of trusted execution environments obtains a public key associated with a subsequent trusted execution environment with which to encrypt any data sent to the subsequent trusted execution environment.

(93) Example 15: A processing system, comprising: a processor (502); and a computer-readable storage medium (504) having computer-executable instructions stored thereupon that, when executed by the processor (502), cause the processing system to: establish, to a third-party attestation service (390), that the processing system (502) is a trusted execution environment; receive a public key (392) and a corresponding private key (394) from the third-party attestation service (390); receive user-specific web usage data (366) generated by a user (102) and that has been encrypted with the public key (392) provided to the first trusted execution environment (343) of a plurality of trusted execution environments by the third-party attestation service (390); decrypt the user-specific web usage data (366) using the private key (394); within the first trusted execution environment (343) of the plurality of trusted execution environments, train a differentially private machine learning model (334) with the user-specific web usage data (366); receive a content request (362) from a client computing device (104); within a second trusted execution environment (341) of the plurality of trusted execution environments: select a piece of content (364) from a plurality of possibly relevant pieces of content (382) to be displayed to the user (102) based on the differentially private machine learning model (334); and transmit the selected content (364) to the client computing device (104) in response to the content request (362).

(94) Example 16: The processing system of Example 15, wherein the differentially private machine learning model is trained in part based on the web usage data.

(95) Example 17: The processing system of Example 15, wherein the user-specific web usage data comprises clicks of content or impressions of content.

(96) Example 18: The processing system of Example 15, wherein the content request originates from an embedded website associated with the processing system.

(97) Example 19: The processing system of Example 15, wherein the instructions further cause the processing system to: incorporate the user-specific web usage data into a differentially private dataset, wherein selecting the piece of content to be displayed to the user is additionally based on the differentially private dataset.

(98) Example 20: The processing system of Example 19, wherein functions applied to the differentially private dataset or the differentially private machine learning model return results that are themselves differentially private.

(99) While certain example embodiments have been described, these embodiments have been presented by way of example only and are not intended to limit the scope of the inventions disclosed herein. Thus, nothing in the foregoing description is intended to imply that any particular feature, characteristic, step, module, or block is necessary or indispensable. Indeed, the novel methods and systems described herein may be embodied in a variety of other forms; furthermore, various omissions, substitutions and changes in the form of the methods and systems described herein may be made without departing from the spirit of the inventions disclosed herein. The accompanying claims and their equivalents are intended to cover such forms or modifications as

would fall within the scope and spirit of certain of the inventions disclosed herein.

(100) It should be appreciated that any reference to “first,” “second,” etc. elements within the Summary and/or Detailed Description is not intended to and should not be construed to necessarily correspond to any reference of “first,” “second,” etc. elements of the claims. Rather, any use of “first” and “second” within the Summary, Detailed Description, and/or claims may be used to distinguish between two different instances of the same element.

(101) In closing, although the various techniques have been described in language specific to structural features and/or methodological acts, it is to be understood that the subject matter defined in the appended representations is not necessarily limited to the specific features or acts described. Rather, the specific features and acts are disclosed as example forms of implementing the claimed subject matter.

Claims

1. A method comprising: establishing, to a third-party attestation service, that a processing system is a trusted execution environment; receiving a public key and a corresponding private key from the third-party attestation service; receiving user-specific web usage data that has been encrypted with the public key; decrypting the user-specific web usage data using the private key; storing the user-specific web usage data in a differentially private data store; selecting content to be displayed to a user based on the differentially private data store; and transmitting the selected content to a client computing device.
2. The method of claim 1, wherein the differentially private data store comprises a differentially private machine learning model or a differentially private data table.
3. The method of claim 2, further comprising: providing the differentially private machine learning model or the differentially private data table to an external service.
4. The method of claim 2, further comprising: providing the differentially private machine learning model or the differentially private data table to an internal service that selects content for users; and deleting the non-differentially private user-specific web usage data.
5. The method of claim 1, further comprising: performing a k-anonymity check on the content to be displayed to remove any identifying information within the content or within a reference to the content.
6. The method of claim 1, further comprising: providing the client computing device with the public key received from the third-party attestation service.
7. The method of claim 2, wherein a publicly disclosed program executing on the trusted execution environment trains the differentially private machine learning model.
8. A computer-readable storage medium having computer-executable instructions stored thereupon that, when executed by a processing system, cause the processing system to: establish, to a third-party attestation service, that the processing system is a trusted execution environment; receive a public key and a corresponding private key from the third-party attestation service; receive user-specific web usage data generated by a user and that has been encrypted with the public key; decrypt the user-specific web usage data using the private key of the trusted third-party attestation service; within the trusted execution environment, train a differentially private machine learning model with the user-specific web usage data; receive a content request from a client computing device; select content to be displayed to the user based on the differentially private machine learning model; encrypt the selected content; and transmit the encrypted selected content to the client computing device in response to the content request.
9. The computer-readable storage medium of claim 8, wherein individual user data is stored on the client computing device before being processed by the trusted execution environment.
10. The computer-readable storage medium of claim 8, wherein the content is selected from a content store that encrypts the content.

11. The computer-readable storage medium of claim 8, wherein the trusted-execution environment establishes that the processing system is a trusted execution environment based on a private key that is permanently stored in the processor of the processing system.
 12. The computer-readable storage medium of claim 8, further comprising: establishing an attested channel that is used to communicate the user-specific web usage data to the trusted execution environment.
 13. The computer-readable storage medium of claim 8, wherein the trusted execution environment is one of a plurality of trusted execution environments, each having a unique public key/private key pair received from the trusted third-party attestation service, and wherein the plurality of trusted execution environments perform a plurality of operations that collectively select the content to be displayed to the user.
 14. The computer-readable storage medium of claim 13, wherein the plurality of trusted execution environments comprise a sequence of trusted execution environments, and wherein each of the sequence of trusted execution environments obtains a public key associated with a subsequent trusted execution environment with which to encrypt any data sent to the subsequent trusted execution environment.
 15. A processing system, comprising: a processor; and a computer-readable storage medium having computer-executable instructions stored thereupon that, when executed by the processor, cause the processing system to: establish, to a third-party attestation service, that the processing system is a first trusted execution environment of a plurality of trusted execution environments; receive a public key and a corresponding private key from the third-party attestation service; receive user-specific web usage data generated by a user and that has been encrypted with the public key provided to the first trusted execution environment of a plurality of trusted execution environments by the third-party attestation service; decrypt the user-specific web usage data using the private key; within the first trusted execution environment of the plurality of trusted execution environments, train a differentially private machine learning model with the user-specific web usage data; receive a content request from a client computing device; within a second trusted execution environment of the plurality of trusted execution environments: select a piece of content from a plurality of possibly relevant pieces of content to be displayed to the user based on the differentially private machine learning model; and transmit the selected content to the client computing device in response to the content request.
 16. The processing system of claim 15, wherein the differentially private machine learning model is trained in part based on the web usage data.
 17. The processing system of claim 15, wherein the user-specific web usage data comprises clicks of content or impressions of content.
 18. The processing system of claim 15, wherein the content request originates from an embedded website associated with the processing system.
 19. The processing system of claim 15, wherein the instructions further cause the processing system to: incorporate the user-specific web usage data into a differentially private dataset, wherein selecting the piece of content to be displayed to the user is additionally based on the differentially private dataset.
 20. The processing system of claim 19, wherein functions applied to the differentially private dataset or the differentially private machine learning model return results that are themselves differentially private.
-