

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent	12393850
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Padmanabhan; Venkata N. et al.

---

### Inferring latency sensitivity of user activity

---

#### Abstract

The systems and methods may analyze an impact of latency on user activities by leveraging the variation of latency seen in the normal course of user activities with an application. The systems and methods may infer the latency sensitivity of users by comparing a biased latency distribution of user actions to an estimate of the underlying unbiased latency distribution. The systems and methods may compute a normalized latency preference of the users using a biased latency probability distribution function and an unbiased latency probability distribution function. The systems and method may use the normalized latency preference to analyze an impact of latency on the user activities.

---

**Inventors:** Padmanabhan; Venkata N. (Karnataka, IN), Saxena; Rohan (Pittsburgh, PA), Thakkar; Parth Dhaval (Champaign, IL)

**Applicant:** Microsoft Technology Licensing, LLC (Redmond, WA)

**Family ID:** 1000008765712

**Assignee:** Microsoft Technology Licensing, LLC (Redmond, WA)

**Appl. No.:** 17/566279

**Filed:** December 30, 2021

#### Prior Publication Data

<b>Document Identifier</b>	<b>Publication Date</b>
US 20230134206 A1	May. 04, 2023

#### Foreign Application Priority Data

IN	202141050020	Nov. 01, 2021
----	--------------	---------------

---

#### Publication Classification

**Int. Cl.:** G06N5/02 (20230101); G06F9/451 (20180101)

**U.S. Cl.:**

**CPC** G06N5/02 (20130101); G06F9/451 (20180201);

## Field of Classification Search

**CPC:** G06N (5/02); G06N (20/00); G06F (9/451)

---

## References Cited

### U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
11405329	12/2021	Cuthbert	N/A	G06N 20/00
2013/0081001	12/2012	McColl	717/128	G06F 11/3419
2020/0034882	12/2019	Deng	N/A	G06Q 30/0244

### OTHER PUBLICATIONS

Asrese, et al., “Measuring Web Latency and Rendering Performance: Method, Tools, and Longitudinal Dataset”, In Journal of IEEE Transactions on Network and Service Management, vol. 16, Issue 2, Jun. 1, 2019, pp. 535-549. cited by applicant

Bai, et al., “Understanding and Leveraging the Impact of Response Latency on User Behaviour in Web Search”, In Journal of ACM Transactions on Information Systems, vol. 36, Issue 2, Aug. 29, 2017, 42 Pages. cited by applicant

“International Search Report and Written Opinion Issued in PCT Application No. PCT/US22/045830”, Mailed Date: Jan. 23, 2023, 13 Pages. cited by applicant

Invitation to file search results or a statement of non-availability pursuant to Rule 70b(1) Received for European Application No. 22800450.3, mailed on Jan. 17, 2025, 01 Page. cited by applicant

---

*Primary Examiner:* Leroux; Etienne P

*Attorney, Agent or Firm:* RAY QUINNEY & NEBEKER

---

## Background/Summary

### CROSS-REFERENCE TO RELATED APPLICATIONS

(1) This application claims priority to and the benefit of Indian Patent Application No. 202141050020, filed on Nov. 1, 2021, which is hereby incorporated by reference in its entirety.

### BACKGROUND

(2) Users engage with a variety of online services, such as email, search, e-commerce, and more. Latency is a key metric for defining the user experience in the context of online services. Users generally react negatively to latency, such that, the higher the latency of a service, the less activities that users are likely to perform using the service.

### BRIEF SUMMARY

(3) This Summary is provided to introduce a selection of concepts in a simplified form that are further described below in the Detailed Description. This Summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used as an aid in

determining the scope of the claimed subject matter.

(4) Some implementations relate to a method for identifying latency preferences of users. The method may include determining a biased distribution of latency of a plurality of user actions for an application over a timeframe, wherein the biased distribution of latency includes a latency for each user action of the plurality of user actions. The method may include inferring an unbiased distribution of latency based on the biased distribution of latency by selecting random times within the timeframe of the biased distribution of latency and using the latency for a user action at or close to the chosen random times for the unbiased distribution of latency. The method may include computing a latency preference of the plurality of user actions as a ratio of the probability density function of the biased distribution of latency and the probability density function of the unbiased distribution of latency. The method may include computing a normalized latency preference by dividing the latency preference by the latency preference corresponding to a reference latency. The method outputting the normalized latency preference as a function of latency.

(5) Some implementations relate to a method for determining an unbiased latency. The method may include obtaining a biased distribution of latency that includes a plurality of user actions with an associated latency over a timeframe for an application. The method may include selecting random points in time of the timeframe. The method may include identifying latency samples from the plurality of user actions in the biased distribution of latency based on the points in time selected. The method may include inferring an unbiased distribution of latency for the application using the latency samples.

(6) Additional features and advantages will be set forth in the description that follows. Features and advantages of the disclosure may be realized and obtained by means of the systems and methods that are particularly pointed out in the appended claims. Features of the present disclosure will become more fully apparent from the following description and appended claims, or may be learned by the practice of the disclosed subject matter as set forth hereinafter.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

(1) In order to describe the manner in which the above-recited and other features of the disclosure can be obtained, a more particular description will be rendered by reference to specific implementations thereof which are illustrated in the appended drawings. For better understanding, the like elements have been designated by like reference numbers throughout the various accompanying figures. While some of the drawings may be schematic or exaggerated representations of concepts, at least some of the drawings may be drawn to scale. Understanding that the drawings depict some example implementations, the implementations will be described and explained with additional specificity and detail through the use of the accompanying drawings in which:

(2) FIG. 1 illustrates an example environment for analyzing the latency of application in accordance with implementations of the present disclosure.

(3) FIG. 2 illustrates an example graph illustrating different samples of user activities over time and the measured latency for the user activities in accordance with implementations of the present disclosure.

(4) FIG. 3 illustrates an example graph illustrating probability density functions (PDF) of the biased distribution of latency and the unbiased distribution of latency in accordance with implementations of the present disclosure.

(5) FIG. 4 illustrates an example graph illustrating a latency preference in accordance with implementations of the present disclosure.

(6) FIG. 5 illustrates an example graph illustrating a normalized latency preference for different

user actions for an application in accordance with implementations of the present disclosure.

(7) FIG. 6 illustrates an example graph illustrating a normalized latency preference for a user action for different groups of users in accordance with implementations of the present disclosure.

(8) FIG. 7 illustrates an example graph illustrating a normalized latency preference for different groups of users for a user action in accordance with implementations of the present disclosure.

(9) FIG. 8 illustrates an example graph illustrating a normalized latency preference for different user actions across different months of the year in accordance with implementations of the present disclosure.

(10) FIG. 9 illustrates an example graph illustrating a normalized latency preference for a user action for different times of day in accordance with implementations of the present disclosure.

(11) FIG. 10 illustrates an example graph illustrating a time-based activity confounding factor for a user action in accordance with implementations of the present disclosure.

(12) FIG. 11 illustrates an example table illustrating adjusting a normalized latency preference based on a time-based activity confounding factor in accordance with implementations of the present disclosure.

(13) FIG. 12 illustrates an example method for identifying latency preferences in accordance with implementations of the present disclosure.

(14) FIG. 13 illustrates an example method for determining an unbiased latency in accordance with implementations of the present disclosure.

#### DETAILED DESCRIPTION

(15) This disclosure generally relates to analyzing the latency of application services. Latency is a key metric for defining the user experience in the context of online services, such as, email, search, and e-commerce. Users generally react negatively to latency. The higher the latency of a service, the less activities that users are likely to perform using the service, which in turn means lost revenue for commercial services. Currently, the analysis of the impact of latency on online services for user activity is based on active intervention (e.g., deliberately manipulating or degrading the latency for some users some of the time and then measuring the impact on user activity).

(16) The present disclosure automatically infers the impact of latency of services on user activities based on analyzing natural experiments (e.g., leveraging the variation in user-experienced latency seen in the normal course of user activity). The present disclosure compares the distribution of latency of the user actions actually performed by the users (e.g., a biased distribution of latency, which reflects the impact, if any of latency on user activity) when interacting with the service with the underlying distribution of latency independent of whether the users choose to perform any action (e.g., an unbiased distribution of latency, which reflects the underlying latency of actions without regard to when user actions are actually performed). The unbiased distribution of latency is estimated by sampling random points in time and picking out the temporally nearest samples from the biased distribution. If the biased distribution is shifted to the left (towards lower latency) as compared to the unbiased distribution, the shift may indicate a greater likelihood of user activity when the latency is lower.

(17) The comparison of the biased distribution of latency and the unbiased distribution of latency helps compute the user preference based on latency. Specifically, with the biased and unbiased distributions expressed as Probability Density Functions (PDFs), computing the ratio of the biased PDF to the unbiased PDF yields the latency preference. Furthermore, to facilitate comparison of the latency preference across user groups, user actions, etc., the latency preference is normalized with respect to a chosen reference latency, yielding a normalized latency preference, which quantifies the relative likelihood of user activity at different levels of latency (the normalized latency preference would be equal to 1 at the reference latency, and greater or lower than 1 at latencies that are more or less preferred by users).

(18) The present disclosure may optionally account for confounding factors, such as, time of day effects, content driven user activity preferences, user conditioning based on the typical latency they

experience and so have come to expect, and/or different classes of users (business customers that pay for the service versus users that use the service for free). Such confounding factors could each impact the level of user activity, separate from the impact of latency. The present disclosure may adjust the normalized latency preference based on the confounding factors to reduce or minimize the effect of the confounding factors on the normalized latency preference.

(19) One technical benefit of the present disclosure is helping guide aimed at the improvement of latency of an application. The analysis of the present disclosure may, for instance, help identify individual actions within an application as the ones that are most latency sensitive and so deserve the greatest attention. For example, the analysis may focus on the inferred latency preference of user actions for selecting a mail item, switching folders, and/or searching for an online email service. The present disclosure may use the analysis to enable prioritization of areas of the application to modify for change where improvement in latency of the application may make the biggest impact on a performance of the application and/or the users actions using the application.

(20) By leveraging the variation in latency that happens in the normal course of user actions, the present disclosure uses a passive approach for analyzing the impact of latency on user activities, and thus, totally avoids the risk of impact on the user activities with active intervention.

(21) Referring now to FIG. 1, illustrated is an example environment **100** for analyzing the latency of applications. The environment **100** may have a plurality of users **104** (e.g., **104a** to **104n**, where  $n$  is a positive integer) interacting with a plurality of devices **102** (e.g., **102a** to **102n**) to access one or more applications **10**. The users **104** may be located in different locations worldwide. In some implementations, the applications **10** may be used to access services provided by the server **106**. Examples of the applications **10** include, but are not limited to, email applications, an online shopping application, and/or a search application. When a user **104** initiates a user action **12** for the application **10** (e.g., opening an email by clicking on the email, clicking a button to place a product in a shopping cart, initiating a search by pressing a search button), the server **106** receives the user action **12** and provides a response **14** (e.g., opening the email, placing the product in the shopping cart, performing a search and providing search results) to the application **10**.

(22) A latency log **16** may record a measured latency **20** for different user actions **12** for the application **10**. The latency **20** is a time it takes for the user **104** to get a response **14** for a user action **12**. For example, the latency **20** may be the measured time from when the user initiated the user action **12** (e.g., clicking on an email, initiating a search by pressing the search button) until the response **14** is received by the user **104**. As such, the latency **20** may be indicative of the user's **104** experience with the user action **12**. A higher latency may result in a slower experience for the user **104** with the user action **12** and a lower latency may result in a faster experience for the user **104** with the user action **12**. Thus, if the latency **20** is large (e.g., the total time from the start of an action until a response is received), the latency **20** may impact the user's **104** willingness to continue using the application **10** and/or performing the user action **12**.

(23) The latency log **16** includes the user action **12** (e.g., selecting a mail item, switching between mail folders) and a time **18** the user action **12** started. The time **18** may be based on timestamps recorded at the server **106** for the user action **12**. The latency log **16** also includes a measured end-to-end latency **20** of the user action **12**. The latency **20** may be measured by the client (e.g., web browser) as the time interval from when the user **104** initiates an action (e.g., clicking a mail item) until the end of the user action **12** (e.g., when the mail item is rendered on the device **102**). The measured latency **20** may be measured by the devices **102** and conveyed to the server **106**, where the user action **12**, the time **18** the user action **12** is initiated by the user **104**, and the measured latency **20** are logged in the latency log **16** for the application **10**. The latency **20** may also be measured by the server **106**.

(24) The latency log **16** may also include metadata **22**. The metadata **22** may provide additional information about the user **104** performing the user action **12** and/or provide additional information about the user action **12**. The metadata **22** may be obtained from user profile information of the

users **104** and/or a context of the users **104** (e.g., time of day, location of the user). Example metadata **22** includes, but is not limited to, a subscription type of the user **104** (e.g., whether the user **104** is a business user paying for the service or a consumer user using the service for free), a type of the user action **12**, a location of the user **104** when accessing the application **10**, an indication of the quality network connectivity of the user **104**, a time of day the user actions **12** occurred, and/or date information for when the user actions **12** occurred. The metadata **22** may be used to identify characteristics of the users **104** and/or a context of the user action **12** and segregate the analysis of the data included in the latency log **16** based on user groups and/or different contexts.

(25) The latency log **16** includes a tuple of data (the time **18**, the user action **12**, the latency **20**, and/or any metadata **22**) for every user action **12** received from the plurality of users **104** for the application **10**. Each entry in the latency log **16**, is annotated with a type of user action **12** (e.g., placing an item in a shopping cart, initiating a search, opening an email message), the time **18** the user action started (e.g., a timestamp indicating when a user action **12** is initiated), and the measured latency **20** for the user action **12** (e.g., from the time an action is initiated until a response is received). The latency log **16** may also include the metadata **22** for the user **104** (e.g., an anonymized user identification of the user **104**). The latency log **16** includes multiple types of user actions **12** from different users **104** using the application **10**. For example, the latency log **16** includes several billion user actions **12** received from millions of users **104** worldwide for accessing an online web email service. The latency log **16** aggregates the user actions **12** received from each of the devices **102** of the plurality of users **104**.

(26) The identity of the users **104** and the content of the data accessed by the users **104** is of no concern to the data processing module **108**. As such, the latency log **16** maintains the data privacy of the users **104** by applying different techniques to the user actions **12** for abstracting the data for the user actions **12** without identifying the content of the data (e.g., identifying the user actions **12** as opening the mail item without looking at the contents of the mail item opened) and without identifying the users **104**. The latency log **16** may abstract the metadata **22**, for example, from the user profiles of the user or a user identification of the user to identify characteristics of the user **104** (e.g., a business customer, a user located in Europe) without identifying the user **104**.

(27) The server **106** may be in communication with a data processing module **108** that processes the latency logs **16** and analyzes the latency **20** of the different user actions **12**. The data processing module **108** may use one or more machine learning models to process and/or analyze the large volume of data obtained in the latency logs **16**.

(28) The data processing module **108** may determine a biased distribution of latency **26** for the user actions **12** for the application **10** by the plurality of users **104**. To an extent a user **104** has preferences for latency **20**, the biased distribution of latency **26** reflects the user's **104** preferences since the latency logs **16** are for the user actions **12** already taken by the user **104**. If the users **104** dislike high latency, the users **104** may tend to perform fewer user actions **12** for an application when the latency **20** is high as compared to when the latency **20** is low, and thus, exhibiting a bias towards lower latency. For example, when the service for the application **10** is fast and responsive (e.g., a lower latency), the users **104** may be more likely to stay on the application **10** and perform more actions, and if the service of the application **10** is slow and unresponsive (e.g., a higher latency), the users **104** may prefer to take a break from the application **10**. User activity (e.g., user actions **12** for the application **10**) may be concentrated in periods of low latency, and thus, provides a biased view of the underlying latency distribution.

(29) The data processing module **108** may use the latency log **16** to construct the biased distribution of latency **20** for the user actions **12** for the application **10**. The data processing module **108** uses the latency **20** of each user action **12** at the time **18**, as recorded in the latency log **16**, to determine the biased distribution of latency **26** for the user actions **12**.

(30) The data processing module **108** may construct a biased probability density function (PDF) **30**

of the biased distribution of latency **26**. The biased PDF **30** reflects the bias, if any, on the part of users **104** to perform more frequent or less frequent user actions **12** based on the latency **20**. In some cases, such bias may arise because of explicit user preference (e.g., the users **104** may use a service less when the latency **20** is high). In other cases, the latency **20**, being in the users' **104** critical path, may slow the users **104** down, and thereby, result in fewer user actions **12**. For example, a user **104** who is scanning through the new emails that have arrived in an inbox may get slowed down if the action of clicking on and opening each email takes longer.

(31) The data processing module **108** may use the biased distribution of latency **26** to infer an unbiased distribution of latency **28**. The unbiased distribution of latency **28** may be estimated by sampling random points in time and picking out the temporally nearest samples from the biased distribution of latency **26**. The unbiased distribution of latency **28** may reflect the inherent or underlying latency distribution independent of the user actions **12**. The unbiased distribution of latency **28** needs to be inferred by the data processing model **108** through indirect means, since the unbiased distribution of latency **28** corresponds to the latency **20** at times that are unrelated to when the users **104** actually made accesses, and therefore, the data processing module **108** may not have direct latency measurements because the data processing module **108** is using the latency logs **16** from the natural interactions of user activity with the application **10** (e.g., when the user actions **12** naturally occur from the users **104** with the application **10**).

(32) The data processing module **108** may approximate the unbiased distribution of latency **28** samples by repeatedly picking points in time **24** uniformly at random and picking a latency sample from the biased distribution of latency **26** that is closest to the chosen point in time **24**. The data processing module **108** may select the latency sample (e.g., the latency **20** of a user action **12** at a time **18**) that is closest in time to the chosen point in time **24**. In some implementations, if there are multiple latency samples at the chosen point in time **24**, the data processing module **108** may pick one of the latency samples at random (e.g., selecting one of the measured latency **20** of the user actions **12** with times **18** that are close to the chosen point in time **24**). In some implementations, if there are multiple latency samples at the chosen point in time, the data processing module **108** may take an average of the latency samples (e.g., take an average of the measured latency **20** of the user actions **12** with times **18** that are close to the chosen point in time **24**). By taking the latency samples at random times, the data processing module **108** may get a sample of the measured latency **20** at times not influenced by the user's **104** choice.

(33) The data processing module **108** may construct an unbiased probability density function (PDF) **32** of the unbiased distribution of latency **28**. In an implementation, the biased PDF **30** and the unbiased PDF **32** may be computed as histograms with a time bin of 10 milliseconds (ms).

(34) The data processing module **108** may also calculate a latency preference **34** corresponding to each latency **20**. The data processing module **108** calculates the raw latency preference **34** as the ratio of:

$$B/U, \quad (1)$$

where B is the biased PDF **30** and U is the unbiased PDF **32**. The latency preference **34** may be a noisy curve, and thus, the data processing module **108** may perform processing to smooth the latency preference **34**. In some implementations, the data processing module **108** uses the Savitzky-Golay filter, with a window size of 101 and polynomial degree of 3 to smooth the latency preference **34** estimate. Any type of processing may be performed by the data processing module **108** to smooth the latency preference **34**.

(35) The data processing module **108** may also select a preference corresponding to a reference latency **36** and may normalize the latency preference **34** to obtain a normalized latency preference **38**. The data processing module **108** may divide the other latency values in the latency preference **34** by the latency preference corresponding to the reference latency **36** to generate the latency values in the normalized latency preference **38**. For example, a normalized latency preference of x (e.g., 0.8) at a particular level of latency means that all other factors (e.g., confounding factors **40**)

being equal, the user **104** is  $(1-x) \times 100\%$  (e.g., 20%) less active at this latency as compared to the reference latency **36**. As such, the normalized latency preference **38** may be equal to 1 at the reference latency **36**, and greater or lower than 1 at latencies that are more or less preferred by the users **104**.

(36) The data processing module **108** may also identify one or more confounding factors **40** that may impact the user actions **12** with the application **10**. The data processing module **108** may adjust the normalized latency preference **38** based on the confounding factors **40** to mitigate or minimize an effect of the confounding factors **40** on the normalized latency preference **38**. The confounding factors may impact the level of user activity, separate from the impact of latency.

(37) One example confounding factor **40** includes the time of day or the day of week when the user actions **12** occurred. The time of day or the day of week may have a significant impact on user activity. For example, the users **104** may be less likely to be active during the middle of the night than during daytime, regardless of the latency. Likewise, the users **104** may be less (or, depending on the service, more) active during the weekend than during the weekdays, again, regardless of the latency. The normalized latency preference **38** may illustrate fewer user actions **12** (e.g., accesses by the users **104** when latency is low) not because the users **104** have an aversion to low latency but because of the time confounder (e.g., users may be less active at night).

(38) The data processing module **108** may adjust the normalized latency preference **38** to account for the time of day or the day of the week confounding factor **40** by pooling together data from across different hours in the day and modeling the time confounder as a time-based activity factor that reflects how active the user **104** is during a particular time of day. For example, the time-based activity factor may likely be high during the daytime and low in the middle of the night. The data processing module **108** may estimate the time-based activity factor using different time slots. The data processing module **108** may compute an average of the time-based activity factor across different latencies for the different time slots to estimate the overall time-based activity confounding factor **40**.

(39) The data processing module **108** may adjust the normalized latency preference **38** based on the estimated time-based activity confounding factor **40**. The adjustment helps neutralize the time confounding factor **40**. For example, the low count of user actions **12** in the middle of the night may be replaced with a higher count of user actions **12** commensurate with a greater prevalence of low latency during the nighttime.

(40) Another example confounding factor **40** includes content driven user activity preferences. Depending on the content, the user **104** may be more or less active, independent of other factors, such as, the latency **20**. The content of the user actions **12** may also impact the volume of activity. The user actions **12** may be of different types, each entailing a different level of user engagement (e.g., clicking on an email, performing a search). In addition, the content returned by the server **106** may determine future actions (e.g., whether the relevant results are returned at the top in which case the user **104** may be led away from a search service, or the results may be a poor match for the search query in which case the user **104** might invoke another search after refining the search terms). The data processing module **108** may adjust the normalized latency preference **38** based on the content driven user activity preference confounding factor **40**, which may help neutralize the content driven user activity preference confounding factor **40** on the latency preference.

(41) Another confounding factor **40** includes user conditioning based on the typical latency the users' experience and have come to expect. User conditioning may have a bearing on the latency sensitivity of the users **104**. For example, if a user **104** expects low latency because the user **104** has a strong network connection, the user **104** may react more negatively than other users who may be accustomed to poor latency (e.g., users in geographical areas with slower network connections). The data processing module **108** may adjust the normalized latency preference **38** based on the user conditioning confounding factor **40** to minimize or mitigate an effect of the user conditioning confounding factor **40** on the normalized latency preference **38**.



(42) The data processing module **108** may analyze the user actions **12** based on different action types and determine a normalized latency preference **38** separately for each action type. The data processing module **108** may also analyze the user actions **12** based on different user groups and determine separate normalized latency preferences **38** for each user group. The data processing module **108** may also analyze a combination of different user actions **12** and different user types and generate different normalized latency preferences **38** for a combination of user groups and user action types. As such, the data processing module **108** may generate the normalized latency preferences **38** for different action types and/or user groups. The normalized latency preference **38** may be output and used by individuals to identify an impact of latency **20** on the user activities (e.g., the user actions **12** for the application **10**).

(43) The data processing module **108** may generate one or more recommendations **42** based on the normalized latency preference **38**. The one or more recommendations **42** may automatically identify areas (e.g., features or functionalities) of the application **10** that may have an impact of latency **20** based on the normalized latency preference **38**. Individuals may use the recommendations **42** to prioritize the identified areas of the application **10** for latency improvement.

(44) For example, the recommendations **42** may provide an insight that the select mail feature of a mail application is more sensitive to latency (i.e., has a steeper latency preference curve) as compared to a compose and send feature of the mail application. Thus, the service owner of the application **10** may focus their efforts based on the insights provided in the recommendation **42** to reduce the latency of the select mail feature, and thus, increase the level of user activity **12**.

(45) The environment **100** may have multiple machine learning models running simultaneously. In some implementations, one or more computing devices (e.g., the server **106** and/or the devices **102**) are used to perform the processing of environment **100**. The one or more computing devices may include, but are not limited to, server devices, personal computers, a mobile device, such as, a mobile telephone, a smartphone, a PDA, a tablet, or a laptop, and/or a non-mobile device. The features and functionalities discussed herein in connection with the various systems may be implemented on one computing device or across multiple computing devices. For example, the applications **10**, the latency logs **16**, and/or the data processing module **108** are implemented wholly on the same computing device. Another example includes one or more subcomponents of the applications **10**, the latency logs **16**, and/or the data processing module **108** implemented across multiple computing devices. Moreover, in some implementations, the applications **10**, the latency logs **16**, and/or the data processing module **108** are implemented or processed on different server devices of the same or different cloud computing networks. Moreover, in some implementations, the features and functionalities are implemented or processed on different server devices of the same or different cloud computing networks.

(46) In some implementations, each of the components of the environment **100** is in communication with each other using any suitable communication technologies. In addition, while the components of the environment **100** are shown to be separate, any of the components or subcomponents may be combined into fewer components, such as into a single component, or divided into more components as may serve a particular embodiment. In some implementations, the components of the environment **100** include hardware, software, or both. For example, the components of the environment **100** may include one or more instructions stored on a computer-readable storage medium and executable by processors of one or more computing devices. When executed by the one or more processors, the computer-executable instructions of one or more computing devices can perform one or more methods described herein. In some implementations, the components of the environment **100** include hardware, such as a special purpose processing device to perform a certain function or group of functions. In some implementations, the components of the environment **100** include a combination of computer-executable instructions and hardware.

(47) As such, the environment **100** may be used to automatically analyze an impact of latency on user activities (e.g., the user actions **12**) of the applications **10**. The environment **100** leverages the variation in latency that happens in the normal course of user actions **12** with the application **10** without active intervention (e.g., deliberately manipulating the latency for some users and measuring the impact on user activity).

(48) Referring now to FIG. 2, illustrated is a graph **200** illustrating different samples of user activities (e.g., user actions **12** with an application **10**) over time and the measured latency for the user activities. The y-axis **202** illustrates the measurement of latency in milliseconds (ms) and the x-axis **204** illustrates time. The measured latency may be determined based on the latency **20** (FIG. 1) included in the latency logs **16** (FIG. 1) for the user action **12** (FIG. 1). For each user action **12** included in the latency logs **16**, the latency logs **16** indicate the time **18** the user action **12** occurred and the measured latency **20** for the user action **12**. The graph **200** illustrates the raw datapoints for the user activities **206** received from the latency logs **16**. The raw datapoints are used to approximate the unbiased distribution of latency **28** (FIG. 1).

(49) The graph **200** illustrates the biased distribution of the user activities **206**, the distribution of latency over time for the user actions **12**. At a certain point in time on the x-axis **204**, the user **104** performed a user action **12**, as indicated by the time **18** recorded in the latency log **16**, and the graph **200** illustrates the experienced latency for the sampled user activities **206** at the time **18**. For example, a timeframe may be selected (e.g., two months) for the x-axis **204** and the user actions **12** that occurred within the timeframe may be presented on the graph **200**.

(50) To an extent a user has preferences for latency, the biased distribution already reflects the user's preferences. For example, if the user **104** prefers lower latency, more users' activities **206** occur during the lower latencies (e.g., the user **104** performed more user actions **12**) and less user activities **206** occur during the higher latencies (e.g., the user **104** performed less user actions **12**), as illustrated in the graph **200**.

(51) The graph **200** also illustrates sampled datapoints **208** within the user activities **206** that are used for constructing the unbiased distribution of latency **28** (FIG. 1). The sampled datapoints **211**, **214**, **218**, **222**, **226**, **230**, **234**, **238** are selected at random times **210**, **212**, **216**, **220**, **224**, **228**, **232**, **236** on the x-axis **204**. In some implementations, the sampled datapoints **211**, **214**, **218**, **222**, **226**, **230**, **234**, **238** are selected as the user activity **12** that is closest in time **18** to the selected times **210**, **212**, **216**, **220**, **224**, **228**, **232**, **236**. For example, for selected time **216**, the sampled datapoint **218** is selected as the user activity **12** that occurred at a time **18** that is closest to the selected time **216** (e.g., the user activity occurred at the same time **18** as the selected time **216** or within a few seconds of the selected time **216**).

(52) In some implementations, the sampled datapoints **211**, **214**, **218**, **222**, **226**, **230**, **234**, **238** are an average of the user activities that are close in time to the selected time **210**, **212**, **216**, **220**, **224**, **228**, **232**, **236**. For example, for the selected time **220**, the sampled datapoint **222** is an average of the user activities that occur close in time to the selected time **220**. A plurality of user activities (e.g., user actions **12**) may have occurred within a timeframe of the selected time **220** (e.g., within five minutes of the selected time **220**) and the average latencies **20** of the plurality of user activities may be used as the sampled datapoint **222** for the selected time **220**.

(53) In some implementations, the sampled datapoints **211**, **214**, **218**, **222**, **226**, **230**, **234**, **238** are randomly selected from the user activities that are close in time to the selected time **210**, **212**, **216**, **220**, **224**, **228**, **232**, **236**. For example, for the selected time **228**, the sampled datapoint **230** is randomly selected from among the different user activities that occurred close in time to the selected time **228**.

(54) As such, the sampled datapoints **211**, **214**, **218**, **222**, **226**, **230**, **234**, **238** may be an approximation of the latency at the chosen random times **210**, **212**, **216**, **220**, **224**, **228**, **232**, **236**, and the data processing module **108** may approximate the unbiased distribution of latency **28** by drawing samples of the user's activities at random from the biased samples corresponding to the

actual user activity **206**.

(55) Referring now to FIG. 3, the graph **300** illustrates the biased probability density function **306** of the biased distribution of latency **26** (FIG. 1) and the unbiased probability density function **308** of the unbiased distribution of latency **28** (FIG. 1). The probability density functions illustrate a distribution of the number of user actions **12** with a latency. The y-axis **302** is the probability density function and the x-axis **304** is the latency **20** (FIG. 1). A higher point on the curve illustrates more samples of user activities (e.g., user actions **12**) at the latency **20** and a lower point on the curve illustrates less samples of user activities (e.g., user actions **12**) at the latency **20**. As such, for any measured latency **20**, the graph **300** illustrates the probability of user activities occurring for the latency **20**. The biased PDF **306** is shifted left slightly compared to the unbiased PDF **308**, and thus, more of the user actions **12** occur towards the left with the lower latencies as compared to the unbiased distribution of the user actions **12**.

(56) Referring now to FIG. 4, the graph **400** illustrates a latency preference **34** (FIG. 1). The latency preference **34** may be computed by the data processing module **108** (FIG. 1) as the raw ratio of B/U and smoothed, where B is the biased distribution of latency **26** and U is the unbiased distribution of latency **28**. The y-axis **402** is the latency preference **34** and the x-axis **404** is the latency **20**.

(57) The raw latency preference curve **406** illustrates, for any value of latency on the x-axis **404**, the latency preference **34** for the latency **20** (e.g., the ratio of dividing the biased distribution of latency by the unbiased distribution of latency). The graph **400** also illustrates a smoothed preference curve **408**. For example, the data processing module **108** may perform additional processing to smooth out the raw latency preference curve **406**. If the user **104** prefers lower latency, the latency preference **34** is higher for lower latencies and lower for higher latencies. For example, in the latency range of 0 to 500 ms, the graph **400** illustrates a disproportionate share of user activities (e.g., user actions **12**) and in the latency range of 2,000 to 2,500 ms, the graph **400** illustrates a lower share of user activities (e.g., user actions **12**).

(58) Referring now to FIG. 5, the graph **500** illustrates a normalized latency preference **38** (FIG. 1) for different user actions **12** (FIG. 1) for an application **10** (FIG. 1) as a function of latency **20** (FIG. 1). The data processing module **108** (FIG. 1) may calculate the normalized latency preference **38** by dividing the latency preference **34** (FIG. 1) by a reference latency **36** (FIG. 1), a selected value of latency **506**. The y-axis **502** is the normalized latency preference **38** and the x-axis **504** is the latency **20**. The selected value of latency **506** is 300 ms. As such, all of the values of the latency preference **34** are divided by the selected value of latency **506**, 300 ms. By normalizing the latency, the graph **500** may compare different normalized latency preference curves for different user actions **12** based on the same reference latency **36**.

(59) An example use case for a web mail service includes different types of user actions **12** for selecting mail, switching folders, searching, and composing and sending a message. The normalized latency preference curve **508** illustrates the normalized latency preference **38** for the selecting mail user action **12**. The normalized latency preference curve **510** illustrates the normalized latency preference **38** for the switching folders user action **12**. The normalized latency preference curve **512** illustrates the normalized latency preference **38** for the searching user action **12**. The normalized latency preference curve **514** illustrates the normalized latency preference **38** for the composing and sending a message user action **12**. For example, at 300 ms, the normalized latency preference **38** is set to 1 for all the different curves, and thus, a common reference point for the curves **508**, **510**, **512**, **514** for the different user actions **12** of an application **10** (e.g., select mail, switch folders, search, compose and send) is determined.

(60) One benefit of the normalized latency preference is the graph **500** may easily show the relative drop in latency preference as compared to the reference latency **506** (e.g., the latency preference dropped **40** percent from the reference latency). For example, the graph **500** illustrates that the normalized latency preference **38** drops sharply for the select mail user action **12** and for the switch

folder user action **12**, reflecting an expectation of “instantaneous” response the users **104** may have for the select mail and switch folder user actions **12**. As the latency grows to 500 ms, 1000 ms, and 1500 ms, respectively, the normalized latency preference **38** drops to 0.88, 0.68, and 0.61, respectively, and thus, indicating that the increase in latency to 500 ms, 1000 ms, and 1500 ms reduces the incidence of user activity by 12%, 32%, and 39%, respectively, relative to the reference latency of 300 ms.

(61) The graph **500** illustrates that the search user action **12** has a less steep drop off in the normalized latency preference **38**, suggesting that users **104** may be conditioned to tolerating a higher latency **20** for the search operation. The compose and send user action **12** is an asynchronous operation, wherein the user interface returns control to the user **104** even as the email is queued up and sent in the background by the server **106**. The asynchronous operation may explain why the normalized latency preference **38** remains nearly flat for the compose and send user action **12**, indicating that there is little sensitivity to latency **20** for the compose and send user action **12**.

(62) The differences in the normalized latency preferences **38** may be used to provide recommendations **42** for identifying different areas of the application **10** to improve the latency of the application **10** and/or the user's **104** experience with the application **10**.

(63) Referring now to FIG. **6**, illustrated is a graph **600** illustrating a normalized latency preference for a user action **12** (FIG. **1**) for different groups of users **104** (FIG. **1**). The y-axis **602** is the normalized latency preference **38** and the x-axis **604** is the latency **20**. For example, the user action **12** is a select mail action for a web mail service and the groups of users **104** are business users who pay for the mail service and consumers who receive the mail service for free. The normalized latency preference curve **606** illustrates the normalized latency preference **38** for the business users for the select mail user action **12**, and the normalized latency preference curve **608** illustrates the normalized latency preference **38** for the consumers for the select mail user action **12**. The graph **600** illustrates a sharper drop in normalized latency preference **38** for the business users as compared to the consumers who receive the mail service for free. The graph **600** may be used to provide recommendations **42** based on insights provided by the latency preference among different user groups. For example, the business users who pay for the service may have less tolerance for latency as compared to the consumers who receive the service for free.

(64) Referring now to FIG. **7**, illustrated is a graph **700** illustrating a normalized latency preference **38** (FIG. **1**) for a user action **12** (FIG. **1**) for different groups of users **104** (FIG. **1**). The y-axis **702** is the normalized latency preference **38** and the x-axis **704** is the latency **20**. The different groups of users may be segmented into quartiles based on a median latency calculated using the network speeds of the users **104**. For example, quartile **1** corresponds to the users with the lowest latency (e.g., fastest network speeds) and quartile **4** corresponds to the users with the highest latency (e.g., slowest network speeds). The data module **108** may use an anonymized user identifier to compute the per-user median latency, which enables grouping the users **104** into quartiles while maintaining the data privacy of the users **104**.

(65) The normalized latency preference curve **706** illustrates the normalized latency preference **38** for the quartile **1** user group. The normalized latency preference curve **708** illustrates the normalized latency preference **38** for the quartile **2** user group. The normalized latency preference curve **710** illustrates the normalized latency preference **38** for the quartile **3** user group. The normalized latency preference curve **712** illustrates the normalized latency preference **38** for the quartile **4** user group.

(66) The graph **700** illustrates a consistent trend, with the sensitivity to latency decreasing progressively from the quartile **1** user group to the quartile **4** user group for the same latency value. The graph **700** may be used to provide insights provided by the latency preference among different user groups. For example, the users **104** who are used to a lower latency may be more sensitive to latency as compared to the users **104** who are used to higher latency.

(67) Referring now to FIG. 8, illustrated is a graph **800** illustrating a normalized latency preference **38** (FIG. 1) for different user actions **12** (FIG. 1) in different months of the year. The y-axis **802** is the normalized latency preference **38** and the x-axis **804** is the latency **20**. For example, the graph **800** illustrates the select mail user action **12** for an online web email service and the switch folder user action **12** for the online web email service. The normalized latency preference curve **806** illustrates the normalized latency preference **38** for a plurality of users for the select mail user action **12** in the month of February. The normalized latency preference curve **808** illustrates the normalized latency preference **38** for a plurality of users for the select mail user action **12** in the month of January. The normalized latency preference curve **810** illustrates the normalized latency preference **38** for a plurality of users for the switch folder user action **12** in the month of February. The normalized latency preference curve **812** illustrates the normalized latency preference **38** for a plurality of users for the switch folder user action **12** in the month of January. The graph **800** may be used to provide insights based on the latency preference among the plurality of users **104**. For example, the graph **800** illustrates consistency in the drop off in the normalized latency preference **38** across the different months, such may suggest that the users' **104** sensitivity to latency **20** for the selected user actions **12** remains stable over the timeframe considered (e.g., the months of January and February).

(68) Referring now to FIG. 9, illustrated is a graph **900** illustrating a normalized latency preference **38** (FIG. 1) for a user action **12** (FIG. 1) for different times of day. The y-axis **902** is the normalized latency preference **38** and the x-axis **904** is the latency **20**. For example, the user action **12** is a select mail user action for an online email service and the different times of day may be split into six hour periods. The normalized latency preference curve **906** illustrates the normalized latency preference **38** for the plurality of users **104** during the time period of 8 am to 2 pm for the select mail user action **12**. The normalized latency preference curve **908** illustrates the normalized latency preference **38** for the plurality of users **104** during the time period of 2 pm to 8 pm for the select mail user action **12**. The normalized latency preference curve **910** illustrates the normalized latency preference **38** for the plurality of users **104** during the time period of 8 pm to 2 am for the select mail user action **12**. The normalized latency preference curve **912** illustrates the normalized latency preference **38** for the plurality of users **104** during the time period of 2 am to 8 am for the select mail user action **12**.

(69) The graph **900** illustrates in each time period, a consistent trend with the normalized latency preference **38** decreasing as latency increases. The graph **900** also illustrates that the drop in the normalized latency preference **38** is sharper during the daytime periods as compared to during the nighttime periods. The graph **900** may be used to provide recommendations **42** based on insights provided by the latency preference among different users. For example, the users **104** that are active late in the night may have a compelling reason to do so, and therefore, may be less sensitive to latency **20**.

(70) Referring now to FIG. 10, illustrated is a graph **1000** illustrating a time-based activity confounding factor **40** (FIG. 1) for a user action **12** (FIG. 1). The graph **1000** illustrates the effect of the time-based activity confounding factor **40** on the normalized latency preference **38** for the user action **12**. The y-axis **1002** is the time-based activity confounding factor **40** and the x-axis **1004** is the latency **20**.

(71) For example, the user action **12** is a select mail user action for an online email service and the different times of day may be split into six hour periods. The normalized latency preference curve **1006** illustrates the normalized latency preference **38** for the plurality of users **104** during the time period of 8 am to 2 pm for the select mail user action **12** adjusted based on the time-based activity confounding factor **40**. The normalized latency preference curve **1008** illustrates the normalized latency preference **38** for the plurality of users **104** during the time period of 2 pm to 8 pm for the select mail user action **12** adjusted based on the time-based activity confounding factor **40**. The normalized latency preference curve **1010** illustrates the normalized latency preference **38** for the

plurality of users **104** during the time period of 8 pm to 2 am for the select mail user action **12** adjusted based on the time-based activity confounding factor **40**. The normalized latency preference curve **1012** illustrates the normalized latency preference **38** for the plurality of users **104** during the time period of 2 am to 8 am for the select mail user action **12** adjusted based on the time-based activity confounding factor **40**. For example, the data processing module **108** may adjust the normalized latency preference **38** based on the calculated time-based activity confounding factor **40**.

(72) The graph **1000** may be used to provide recommendations **42** based on insights provided by the latency preference among different users **104**. For example, the time-based activity confounding factor **40** is lower during the late periods, reflecting a lower level of user activity at night, regardless of the latency **20**. Moreover, the time-based activity confounding factor **40** remains flat across the latency range. The insights may be used to improve the latency **20** of the user action **12**, and thus, improve the users' **104** interaction with the web email service.

(73) Referring now to FIG. **11**, the table **1100** illustrates an adjustment of the normalized latency preference **38** (FIG. **1**) based on a time-based activity confounding factor **40** (FIG. **1**). In the table **1100**, time is discretized into two equal-length slots (“day” and “night”) and latency is divided into two bins (“low latency” and “high latency”). The column **1102** illustrates the time slots (e.g., day or night), the column **1104** illustrates the latency (e.g., “low” or “high”), the column **1106** illustrates a number of user actions (e.g., user actions **12**), the column **1108** illustrates a percentage of the time with this latency, and the column **1110** illustrates the normalized number of user actions.

(74) If the data processing module **108** had ignored the time-based activity confounding factor **40**, the data processing module **108** may have computed the user's **104** level of activity when the latency **20** is “low” as

$$(90+24)/(30+80)=1.04 \quad (2)$$

actions per unit time, and that when the latency **20** is “high” as

$$(140+4)/(70+20)=1.6 \quad (3)$$

actions per unit time. The calculation of the user's **104** level of activity indicates that the user **104** performs more actions when the latency **20** is “high” as compared to when the latency **20** is “low.”

(75) Instead, if the data processing module **108** treats the “day” time slot as the reference and normalizes the counts corresponding to the “night” time slot, the time-based factor would be estimated as

$$\alpha_{\text{sub.night,low}}=(26/80)/(90/30)=0.108 \quad (4)$$

and

$$\alpha_{\text{sub.night,high}}=(4/20)/(140/70)=0.100, \quad (5)$$

so

$$\alpha_{\text{sub.night}}=(0.108+0.100)/2=0.104 \quad (6)$$

(e.g., the average of  $\alpha$  across the latency bins). Therefore, the normalized count of actions during the night is

$$26/0.104=250 \text{ and } 4/0.104=38, \quad (7)$$

respectively, for the “low” and “high” latency bins. Combining the normalized counts with those from the day, the user's level of activity is estimated as

$$(90+250)/(30+80)=3.09 \quad (8)$$

actions per unit time when the latency is “low” and as

$$(140+38)/(70+20)=1.97 \quad (9)$$

actions per unit time when the latency is “high.” That is, the level of user activity is higher when the latency is “low” (e.g., more user actions **12**) compared to when the latency is “high” (e.g., fewer user actions **12**). As such, table **1100** illustrates the changes to the normalized latency preference **38** by adjusting for the time-based activity confounding factor **40**.

(76) Referring now to FIG. **12**, illustrated is an example method **1200** for identifying latency preferences of users. The actions of the method **1200** are discussed below with reference to the

architecture of FIG. 1.

(77) At **1202**, the method **1200** includes determining a biased distribution of latency of a plurality of user actions for an application over a timeframe. The data processing module **108** may determine a biased distribution of latency **26** of a plurality of user actions **12** for an application **10**. The biased distribution of latency **26** includes a latency **20** for each user action **12**. The latency **20** of the user actions **12** may be obtained from a latency log **16** on the server **106**. The latency log **16** may aggregate the plurality of user actions **12** received from a plurality of users **104**. The latency log **16** also indicates a time **18** when the user action occurred (e.g., a time at which the user action **12** is initiated by the user **104**).

(78) The latency log **16** may include a tuple of data (the time **18**, the user action **12**, the latency **20**, and/or any metadata **22**) for every user action **12** received from the plurality of users **104** for the application. Each entry in the latency log **16**, is annotated with a type of user action **12** (e.g., placing an item in a shopping cart, initiating a search, opening an email message), the time **18** the user action started (e.g., a timestamp indicating when a user action **12** is initiated), and the measured latency **20** for the user action **12** (e.g., from the time an action is initiated until a response is received).

(79) The latency log **16** may also include the metadata **22** for the user **104** (e.g., an anonymized user identification of the user **104**). The metadata **22** may be obtained from user profile information of the users **104** and/or a context of the users **104** (e.g., time of day, location of the user). Example metadata **22** includes a subscription type of the user **104** (e.g., whether the user **104** is a business user paying for the service or a consumer user using the service for free), a type of the user action **12**, a location of the user **104** when accessing the application **10**, and/or an indication of the quality network connectivity of the user **104**. The metadata **22** may be used to identify characteristics of the users **104** and/or a context of the user action **12** and segregate the analysis of the data included in the latency log **16** based on user groups and/or different contexts.

(80) The data processing module **108** may use the latency log **16** to construct the biased distribution of latency **20** for the user actions **12** for the application **10**. The data processing module **108** uses the latency **20** of each user action **12** at the time **18**, as recorded in the latency log **16**, to determine the biased distribution of latency **26** for the user actions **12**. The data processing module **108** may use one or more machine learning models to process and/or analyze the large volume of data obtained in the latency logs **16**.

(81) At **1204**, the method **1200** includes inferring an unbiased distribution of latency based on the biased distribution of latency. The data processing module **108** may use the biased distribution of latency **26** to infer an unbiased distribution of latency **28**. The unbiased distribution of latency **28** may be estimated by sampling random points in time and picking out the temporally nearest samples from the biased distribution of latency **26**. The unbiased distribution of latency **28** may reflect the inherent or underlying latency distribution independent of the user actions **12**.

(82) The data processing module **108** may approximate the unbiased distribution of latency **28** samples by selecting random points in time **24** within the timeframe of the biased distribution of latency **26** and using the latency **20** for a user action **12** at the chosen random points in time **24**, or close to the chosen points in time **24**, for the unbiased distribution of latency **28**.

(83) The data processing module **108** may select a latency sample (e.g., the latency **20** of a user action **12** at a time **18**) that is closest in time to the chosen point in time **24**. In some implementations, if there are multiple latency samples at the chosen point in time **24**, the data processing module **108** may pick one of the latency samples at random (e.g., selecting one of the measured latency **20** of the user actions **12** with times **18** that are close to the chosen point in time **24**). In some implementations, if there are multiple latency samples at the chosen point in time, the data processing module **108** may take an average of the latency samples (e.g., take an average of the measured latency **20** of the user actions **12** with times **18** that are close to the chosen point in time **24**). By taking the latency samples at random times, the data processing module **108** may get a

sample of the measured the latency **20** at times not influenced by the user's **104** choice.

(84) At **1206**, the method **1200** includes computing a latency preference of the plurality of user actions. The data processing module **108** may compute the latency preference **34** as a ratio of a PDF of the biased distribution of latency (e.g., biased PDF **30**) and a PDF of the unbiased distribution of latency (e.g., unbiased PDF **32**). The data processing module **108** may construct an unbiased PDF **32** of the unbiased distribution of latency **28**. The data processing module **108** may also calculate a latency preference **34** corresponding to each latency. The data processing module **108** calculates the latency preference **34** as the ratio of the biased PDF **30** and the unbiased PDF **32**. The latency preference **34** may be a noisy curve, and thus, the data processing module **108** may perform processing to smooth the latency preference **34**.

(85) At **1208**, the method **1200** includes computing a normalized latency preference. The data processing module **108** may also select a preference corresponding to a reference latency **36** and may normalize the latency preference **34** to obtain a normalized latency preference **38**. The data processing module **108** may divide the other latency values in the latency preference **34** by the latency preference corresponding to the reference latency **36** to generate the latency values in the normalized latency preference **38**.

(86) The data processing module **108** may generate different normalized latency preferences **34** for different groups of users **104**. In addition, the data processing module **108** may generate different normalized latency preferences **34** for different types of user actions **12**. The data processing module **108** may also generate different normalized latency preferences **34** based on the time of day when the plurality of user actions **12** occurred, the location where the plurality of user actions **12** occurred, and/or a date when the plurality of user actions **12** occurred. The data processing module **108** may use a variety of factors or a combination of factors from the metadata **22** in determining the different normalized latency preferences **34**.

(87) The data processing module **108** may also identify one or more confounding factors **40** that may impact the user actions **12** with the application **10**. The data processing module **108** may adjust the normalized latency preference **38** based on the confounding factors **40** to mitigate or minimize an effect of the confounding factors **40** on the normalized latency preference **38**. Example confounding factors **40** include, but are not limited to, a time-based activity factor, content driven user activity preference, or previous user conditioning based on the typical latency the users' experience and have come to expect. Each confounding factor **40** may impact the level of user activity, separate from the impact of latency.

(88) At **1210**, the method **1200** includes outputting the normalized latency preference as a function latency. The data processing module **108** may output the normalized latency preference **38** as a function of latency. The normalized latency preference **38** may be used to generate one or more recommendations **42** for the application functionality to be prioritized for latency improvement. The recommendations **42** may be based on analyzing the normalized latency preference. For example, the data processing module **108** may analyze the different normalized latency preferences **38** (e.g., a normalized latency preference **38** for different groups of users, a normalized latency preference **38** for different action types of user actions, a normalized latency preference **38** for user actions at different times of the day) and generate one or more recommendations **42** based on the analysis. The recommendations **42** may identify one or more areas of the application to modify or change to reduce an amount of latency for the application.

(89) As such, the method **1200** may be used to automatically analyze an impact of latency **20** on user activities **12** by leveraging the variation of latency seen in the normal course of the user activities **12** with an application **10**.

(90) Referring now to FIG. **13**, illustrated is an example method **1300** for determining an unbiased latency. The actions of the method **1300** are discussed below with reference to the architecture of FIG. **1**.

(91) At **1302**, the method **1300** includes obtaining a biased distribution of latency that includes a



plurality of user actions with an associated latency over a timeframe for an application. The biased distribution of latency **26** is based on logs of user actions from a plurality of users **104** interacting with the application **10**.

(92) At **1304**, the method **1300** includes selecting random points in time of the timeframe of the biased distribution of latency. The data module **108** may repeatedly pick points in time **24** at random within the timeframe of the biased distribution of latency **26**.

(93) At **1306**, the method **1300** includes identifying latency samples from the plurality of user actions in the biased distribution of latency based on the points in time selected. The data processing module **108** may identify a measurement from the plurality of user actions that is closest in time **18** to a selected point in time **24** and using the measurement (e.g., the latency **20**) of the user action **12** that is closest in time **18** as a latency sample for the selected point in time **24**. The data processing module **108** may also identify one or more user actions **12** of the plurality of user actions that are close in time **18** to a selected point in time **24** and take an average latency **20** of the one or more user actions **12** as a latency sample for the selected point in time **24**. The data processing module **108** may also identify one or more user actions **12** of the plurality of user actions that are close in time **18** to a selected point in time **24** and randomly selecting a latency **20** of the one or more user actions **12** as a latency sample for the selected point in time **24**.

(94) At **1308**, the method **1300** includes inferring an unbiased distribution of latency for the application using the latency samples. The data processing module **108** may use the biased distribution of latency **26** to infer an unbiased distribution of latency **28**. The data processing module **108** may use the latency samples from the plurality of user actions in the biased distribution of latency **26** to construct the unbiased distribution of latency **26** for the application **10**.

(95) As such, the method **1300** may be used to infer the unbiased distribution of latency **28** and reflect the inherent or underlying latency distribution independent of the user actions **12**.

(96) As illustrated in the foregoing discussion, the present disclosure utilizes a variety of terms to describe features and advantages of the model evaluation system. Additional detail is now provided regarding the meaning of such terms. For example, as used herein, a “machine learning model” refers to a computer algorithm or model (e.g., a classification model, a binary model, a regression model, a language model, an object detection model) that can be tuned (e.g., trained) based on training input to approximate unknown functions. For example, a machine learning model may refer to a neural network (e.g., a convolutional neural network (CNN), deep neural network (DNN), recurrent neural network (RNN)), or other machine learning algorithm or architecture that learns and approximates complex functions and generates outputs based on a plurality of inputs provided to the machine learning model. As used herein, a “machine learning system” may refer to one or multiple machine learning models that cooperatively generate one or more outputs based on corresponding inputs. For example, a machine learning system may refer to any system architecture having multiple discrete machine learning components that consider different kinds of information or inputs.

(97) The techniques described herein may be implemented in hardware, software, firmware, or any combination thereof, unless specifically described as being implemented in a specific manner. Any features described as modules, components, or the like may also be implemented together in an integrated logic device or separately as discrete but interoperable logic devices. If implemented in software, the techniques may be realized at least in part by a non-transitory processor-readable storage medium comprising instructions that, when executed by at least one processor, perform one or more of the methods described herein. The instructions may be organized into routines, programs, objects, components, data structures, etc., which may perform particular tasks and/or implement particular data types, and which may be combined or distributed as desired in various implementations.

(98) Computer-readable mediums may be any available media that can be accessed by a general purpose or special purpose computer system. Computer-readable mediums that store computer-

executable instructions are non-transitory computer-readable storage media (devices). Computer-readable mediums that carry computer-executable instructions are transmission media. Thus, by way of example, and not limitation, implementations of the disclosure can comprise at least two distinctly different kinds of computer-readable mediums: non-transitory computer-readable storage media (devices) and transmission media.

(99) As used herein, non-transitory computer-readable storage mediums (devices) may include random access memory (RAM), read-only memory (ROM), electrically erasable programmable read-only memory (EEPROM), compact disc read-only memory (CD-ROM), solid state drives (“SSDs”) (e.g., based on RAM), Flash memory, phase-change memory (“PCM”), other types of memory, other optical disk storage, magnetic disk storage or other magnetic storage devices, or any other medium which can be used to store desired program code means in the form of computer-executable instructions or data structures and which can be accessed by a general purpose or special purpose computer.

(100) The steps and/or actions of the methods described herein may be interchanged with one another without departing from the scope of the claims. In other words, unless a specific order of steps or actions is required for proper operation of the method that is being described, the order and/or use of specific steps and/or actions may be modified without departing from the scope of the claims.

(101) The term “determining” encompasses a wide variety of actions and, therefore, “determining” can include calculating, computing, processing, deriving, investigating, looking up (e.g., looking up in a table, a database, a datastore, or another data structure), ascertaining and the like. Also, “determining” can include receiving (e.g., receiving information), accessing (e.g., accessing data in a memory) and the like. Also, “determining” can include resolving, selecting, choosing, establishing, predicting, inferring, and the like.

(102) The articles “a,” “an,” and “the” are intended to mean that there are one or more of the elements in the preceding descriptions. The terms “comprising,” “including,” and “having” are intended to be inclusive and mean that there may be additional elements other than the listed elements. Additionally, it should be understood that references to “one embodiment” or “an embodiment” of the present disclosure are not intended to be interpreted as excluding the existence of additional implementations that also incorporate the recited features. For example, any element described in relation to an embodiment herein may be combinable with any element of any other embodiment described herein. Numbers, percentages, ratios, or other values stated herein are intended to include that value, and also other values that are “about” or “approximately” the stated value, as would be appreciated by one of ordinary skill in the art encompassed by implementations of the present disclosure. A stated value should therefore be interpreted broadly enough to encompass values that are at least close enough to the stated value to perform a desired function or achieve a desired result. The stated values include at least the variation to be expected in a suitable manufacturing or production process, and may include values that are within 5%, within 1%, within 0.1%, or within 0.01% of a stated value.

(103) A person having ordinary skill in the art should realize in view of the present disclosure that equivalent constructions do not depart from the spirit and scope of the present disclosure, and that various changes, substitutions, and alterations may be made to implementations disclosed herein without departing from the spirit and scope of the present disclosure. Equivalent constructions, including functional “means-plus-function” clauses are intended to cover the structures described herein as performing the recited function, including both structural equivalents that operate in the same manner, and equivalent structures that provide the same function. It is the express intention of the applicant not to invoke means-plus-function or other functional claiming for any claim except for those in which the words ‘means for’ appear together with an associated function. Each addition, deletion, and modification to the implementations that falls within the meaning and scope of the claims is to be embraced by the claims.

(104) The present disclosure may be embodied in other specific forms without departing from its spirit or characteristics. The described implementations are to be considered as illustrative and not restrictive. The scope of the disclosure is, therefore, indicated by the appended claims rather than by the foregoing description. Changes that come within the meaning and range of equivalency of the claims are to be embraced within their scope.

## Claims

1. A method for identifying latency preferences of users, comprising: determining a biased distribution of latency of a plurality of user actions for an application over a timeframe, wherein the biased distribution of latency includes a latency for each user action of the plurality of user actions; inferring an unbiased distribution of latency based on the biased distribution of latency by selecting random times within the timeframe of the biased distribution of latency and using the latency for a user action at or close to the chosen random times for the unbiased distribution of latency; computing a latency preference of the plurality of user actions as a ratio of a probability density function of the biased distribution of latency and a probability density function of the unbiased distribution of latency; computing a normalized latency preference by dividing the latency preference by the latency preference corresponding to a reference latency; and outputting the normalized latency preference as a function of latency.
2. The method of claim 1, wherein the latency of the plurality of user actions is obtained from a latency log from a server.
3. The method of claim 2, wherein the latency log aggregates the plurality of user actions from a plurality of users interacting with the application.
4. The method of claim 2, wherein the latency log indicates a time when the user actions occurred.
5. The method of claim 2, wherein the latency log further includes: metadata of a plurality of users or metadata for the plurality of user actions, wherein the metadata is obtained from one or more of user profile information, a context of the plurality of users, or an action type of the plurality of user actions.
6. The method of claim 5, wherein the action type identifies a type of the plurality of user actions and the user profile information identifies different groups of the plurality of users.
7. The method of claim 6, wherein different normalized latency preferences are generated for the different groups of the plurality of users.
8. The method of claim 5, wherein different normalized latency preferences are generated for different action types of the plurality of user actions.
9. The method of claim 5, wherein the context of the plurality of users is used to identify one or more of a location where the plurality of user actions occurred, a network connectivity of the users, or a type of user, and wherein different normalized latency preferences are generated based on one or more of the location where the plurality of user actions occurred, the network connectivity of the users, or the type of user.
10. The method of claim 1, further comprising: adjusting the normalized latency preference based on at least one confounding factor by modifying the normalized latency preference to mitigate an effect of the at least one confounding factor on the normalized latency preference, wherein the at least one confounding factor includes a time-based activity factor, content driven user activity preference, or previous user conditioning.
11. The method of claim 1, further comprising: generating one or more recommendations for the application functionality to be prioritized for latency improvement, based on analyzing the normalized latency preference.
12. The method of claim 11, wherein the one or more recommendations identify one or more areas of the application to modify or change to reduce an amount of latency for the application.
13. The method of claim 11, wherein analyzing the normalized latency preference is based on

analyzing different normalized latency preferences of different groups of users for an action type of the plurality of user actions.

14. The method of claim 11, wherein analyzing the normalized latency preference is based on analyzing different normalized latency preferences for a combination of different groups of users and different action types of the plurality of user actions.

15. The method of claim 10, wherein analyzing the normalized latency preference is based on analyzing the normalized latency preference of a plurality of users for the plurality of user actions during different times of day.

16. A method for determining an unbiased latency, comprising: obtaining a biased distribution of latency that includes a plurality of user actions with an associated latency over a timeframe for an application; selecting random points in time of the timeframe; identifying latency samples from the plurality of user actions in the biased distribution of latency based on the points in time selected; and inferring an unbiased distribution of latency for the application using the latency samples.

17. The method of claim 16, wherein the biased distribution of latency is based on logs of user actions from a plurality of users interacting with the application.

18. The method of claim 16, wherein identifying the latency samples includes: identifying a measurement from the plurality of user actions that is closest in time to a selected point in time and using the latency of a nearest user action as a latency sample for the selected point in time.

19. The method of claim 16, wherein identifying the latency sample includes: identifying one or more user actions of the plurality of user actions that are close in time to a selected point in time and taking an average latency of the one or more user actions as a latency sample for the selected point in time.

20. The method of claim 16, wherein identifying the latency sample includes: identifying one or more user actions of the plurality of user actions that are close in time to a selected point in time and randomly selecting a latency of the one or more user actions as a latency sample for the selected point in time.

---