

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12393861
Kind Code	B1
Date of Patent	August 19, 2025
Inventor(s)	Kulkarni; Chinmay Madhav et al.

Systems and methods for securing untrusted code

Abstract

A computer-implemented method is disclosed that includes operations of receiving a request to store a representation of a machine learning (ML) model in a non-transitory computer-readable medium, validating the representation of the ML model, and storing the representation of the ML model. The representation of the ML model may be a computation graph, and further, the computation graph may be a format for representing ML models through a defined common set of operators and a common file format. Validating the representation of the ML model may include verifying an input schema of the representation of the ML model is compatible with an input schema of the ML model. The method may include further operations of deserializing the representation of the ML model, wherein a successful deserialization indicates that the representation of the ML model is valid and secure.

Inventors: Kulkarni; Chinmay Madhav (Sunnyvale, CA), Ma; Lin (Burnaby, CA), Malekpour; Amir (Vancouver, CA), Rajagopalan; Mohan (Mountain View, CA), Reed; John C. (Saratoga, CA), Sriharsha; Ram (Oakland, CA)

Applicant: SPLUNK Inc. (San Francisco, CA)

Family ID: 1000005166278

Assignee: Cisco Technology, Inc. (San Jose, CA)

Appl. No.: 17/074441

Filed: October 19, 2020

Publication Classification

Int. Cl.: G06F16/21 (20190101); G06F16/215 (20190101); G06F16/901 (20190101); G06F18/21 (20230101); G06N20/00 (20190101)

U.S. Cl.:

CPC **G06N20/00** (20190101); **G06F16/215** (20190101); **G06F16/9024** (20190101);
G06F18/217 (20230101);

Field of Classification Search

CPC: G06N (20/00); G06F (16/215); G06F (16/9027); G06F (18/21)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7937344	12/2010	Baum et al.	N/A	N/A
8112425	12/2011	Baum et al.	N/A	N/A
8751529	12/2013	Zhang et al.	N/A	N/A
8788525	12/2013	Neels et al.	N/A	N/A
9215240	12/2014	Merza et al.	N/A	N/A
9286413	12/2015	Coates et al.	N/A	N/A
10127258	12/2017	Lamas et al.	N/A	N/A
2008/0046869	12/2007	Lambert	717/136	G06F 9/4493
2014/0082071	12/2013	Rexer	709/204	G06F 16/192
2015/0127877	12/2014	Novellini	710/316	H03M 9/00
2016/0253446	12/2015	Chong	716/136	G06F 30/36
2019/0098106	12/2018	Mungel et al.	N/A	N/A
2019/0230090	12/2018	Kathiara	N/A	H04L 9/3247
2020/0097851	12/2019	Alvarez	N/A	G05B 13/0265
2020/0380009	12/2019	Reynolds	N/A	G06F 16/2423
2022/0137988	12/2021	Graham	719/320	G06F 9/44521

OTHER PUBLICATIONS

Splunk Enterprise 8.0.0 Overview, available online, retrieved May 20, 2020 from docs.splunk.com. cited by applicant

Splunk Cloud 8.0.2004 User Manual, available online, retrieved May 20, 2020 from docs.splunk.com. cited by applicant

Splunk Quick Reference Guide, updated 2019, available online at <https://www.splunk.com/pdfs/solution-guides/splunk-quick-reference-guide.pdf>, retrieved May 20, 2020. cited by applicant

Carraso, David, "Exploring Splunk," published by CITO Research, New York, NY, Apr. 2012. cited by applicant

Bitincka, Ledion et al., "Optimizing Data Analysis with a Semi-structured Time Series Database," self-published, first presented at "Workshop on Managing Systems via Log Analysis and Machine Learning Techniques (SLAML)", Vancouver, British Columbia, Oct. 3, 2010. cited by applicant

C. M. Kulkarni, et al. "Systems and Methods for a Unified Analytics Platform," filed Oct. 19, 2020, U.S. Appl. No. 17/073,752 including its prosecution history. cited by applicant

C. M. Kulkarni, et al. "Systems and Methods for Integration of Multiple Programming Languages within a Pipelined Search Query," filed Oct. 19, 2020, U.S. Appl. No. 17/074,280 including its prosecution history. cited by applicant

C. M. Kulkarni, et al. "Systems and Methods for Auto-Deployment of a Machine Learning Component within a Pipelined Search Query," filed Oct. 19, 2020, U.S. Appl. No. 17/074,407 including its prosecution history. cited by applicant

C. M. Kulkarni, et al. "Systems and Methods for Integration of Machine Learning Components within a Pipelined Search Query to Generate a Graphic Visualization," filed Oct. 19, 2020, U.S. Appl. No. 17/074,206 Including its prosecution history. cited by applicant
C. M. Kulkarni, et al. "Systems and Methods for Decoupling Search Processing Language and Machine Learning Analytics from Storage of Accessed Data," filed Oct. 19, 2020, U.S. Appl. No. 17/074,100 including its prosecution history. cited by applicant

Primary Examiner: Salomon; Phenuel S

Attorney, Agent or Firm: Rutan & Tucker, LLP

Background/Summary

RELATED APPLICATIONS

(1) Any and all applications for which a foreign or domestic priority claim is identified in the Application Data Sheet as filed with the present application are incorporated by reference under 37 CFR 1.57 and made a part of this specification.

BACKGROUND

(2) Information technology (IT) environments can include diverse types of data systems that store large amounts of diverse data types generated by numerous devices. For example, a big data ecosystem may include databases such as MySQL and Oracle databases, cloud computing services such as Amazon web services (AWS), and other data systems that store passively or actively generated data, including machine-generated data ("machine data"). The machine data can include log data, performance data, diagnostic data, metrics, tracing data, or any other data that can be analyzed to diagnose equipment performance problems, monitor user interactions, and to derive other insights.

(3) The large amount and diversity of data systems containing large amounts of structured, semi-structured, and unstructured data relevant to any search query can be massive, and continues to grow rapidly. This technological evolution can give rise to various challenges in relation to managing, understanding and effectively utilizing the data. To reduce the potentially vast amount of data that may be generated, some data systems pre-process data based on anticipated data analysis needs. In particular, specified data items may be extracted from the generated data and stored in a data system to facilitate efficient retrieval and analysis of those data items at a later time. At least some of the remainder of the generated data is typically discarded during pre-processing.

(4) However, storing massive quantities of minimally processed or unprocessed data (collectively and individually referred to as "raw data") for later retrieval and analysis is becoming increasingly more feasible as storage capacity becomes more inexpensive and plentiful. In general, storing raw data and performing analysis on that data later can provide greater flexibility because it enables an analyst to analyze all of the generated data instead of only a fraction of it. Although the availability of vastly greater amounts of diverse data on diverse data systems provides opportunities to derive new insights, it also gives rise to technical challenges to search and analyze the data in a performant way.

(5) Furthermore, over the last few years, a data analytic platform compatible across different programming languages and different libraries has been urgently needed by the data science community, as data scientists are relying on a wider selection of data analytic development tools. For example, many data scientists are utilizing their JUPYTER® notebooks as their data analytics development tool. JUPYTER® notebooks are client-based interactive web applications that allow users to create and share codes and visualizations. Given that these notebooks provide an

interactive computing environment that supports multiple (two or more) programming languages, it has become a technical challenge to provide a search capability that efficiently and effectively supports multiple programming languages without unacceptable search latency caused from programming language translations.

(6) Additionally, search queries are becoming more complex given the growing usage of artificial intelligence (AI) in the gathering and processing of data. One AI type is referred to as “machine learning,” which generally involves software configured to build a mathematical model based on sample data, sometimes referred to as “training data,” in order to make predictions or decisions without the software being explicitly programmed to make such predictions or decisions. Data scientists are attempting to use ML models and/or ML streaming algorithms as part of query searches into data repositories, but there is no current framework that allows for ML models and/or ML streaming algorithms to be deployed within any search workflow or ML pipeline independent of programming language or origination. Also, there is no mechanism to formulate portions of ML-based pipelines for use and re-use by data scientists in development of pipelined search queries.

(7) Further, data scientists may desire to develop, train and utilize their own ML models within a search query, especially a pipelined search query that is processing on streaming data. However, this is currently no way to ensure these ML models are secure and consist of trusted code. Although execution of untrusted code may be acceptable in a sandboxed environment where malicious code injected into a ML model cannot access others' data or cause harm, untrusted code cannot be executed in distributed runtime environment in which computing resources shared among multiple tenants. Thus, in order to utilize the advantages of execution in a distributed runtime environment, such as resource sharing and computing efficiencies, all ML models utilized within a search query need to be secured.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) Illustrative examples are described in detail below with reference to the following figures:
- (2) FIG. 1 is a block diagram of an embodiment of a data processing environment.
- (3) FIG. 2 is a flow diagram illustrating an embodiment of a routine implemented by the data intake and query system to process, index, and store data.
- (4) FIG. 3A is a block diagram illustrating an embodiment of machine data received by the data intake and query system.
- (5) FIGS. 3B and 3C are block diagrams illustrating embodiments of various data structures for storing data processed by the data intake and query system.
- (6) FIG. 4A is a flow diagram illustrating an embodiment of a routine implemented by the query system to execute a query.
- (7) FIG. 4B provides a visual representation of the manner in which a pipelined command language or query can operate
- (8) FIG. 4C is a block diagram illustrating an embodiment of a configuration file that includes various extraction rules that can be applied to events.
- (9) FIG. 4D is a block diagram illustrating an example scenario where a common customer identifier is found among log data received from disparate data sources.
- (10) FIG. 5A is a block diagram of an embodiment of a data processing environment including a data science platform.
- (11) FIG. 5B is a detailed block diagram of the data processing environment of FIG. 5A.
- (12) FIG. 5C is an illustration of an operator sequence as a logical plan.
- (13) FIG. 6A is a block diagram illustrating a first set of operations performed within the data processing environment of FIG. 5A.

(14) FIG. 6B is a block diagram illustrating a second set of operations performed within the data processing environment of FIG. 5A.

(15) FIGS. 7A-7B are flow diagrams illustrating an embodiment of securing a third-party code and processing of a SPL query.

(16) FIG. 8 is a block diagram illustrating an embodiment of the model service component illustrated in FIG. 5A.

(17) FIG. 9 is a block diagram illustrating an embodiment of the components associated with the model validation component deployed within the model service component illustrated in FIG. 8.

(18) FIGS. 10A-10C are flow diagrams illustrating an embodiment of securing a machine-learning model for storage on the data science platform of FIG. 5A.

DETAILED DESCRIPTION

(19) Modern data centers and other computing environments can comprise anywhere from a few host computer systems to thousands of systems configured to process data, service requests from remote clients, and perform numerous other computational tasks. During operation, various components within these computing environments often generate significant volumes of machine data. Machine data is any data produced by a machine or component in an information technology (IT) environment and that reflects activity in the IT environment. For example, machine data can be raw machine data that is generated by various components in IT environments, such as servers, sensors, routers, mobile devices, Internet of Things (IoT) devices, etc. Machine data can include system logs, network packet data, sensor data, application program data, error logs, stack traces, system performance data, etc. In general, machine data can also include performance data, diagnostic information, and many other types of data that can be analyzed to diagnose performance problems, monitor user interactions, and to derive other insights.

(20) A number of tools are available to analyze machine data. In order to reduce the size of the potentially vast amount of machine data that may be generated, many of these tools typically pre-process the data based on anticipated data-analysis needs. For example, pre-specified data items may be extracted from the machine data and stored in a database to facilitate efficient retrieval and analysis of those data items at search time. However, the rest of the machine data typically is not saved and is discarded during pre-processing. As storage capacity becomes progressively cheaper and more plentiful, there are fewer incentives to discard these portions of machine data and many reasons to retain more of the data.

(21) This plentiful storage capacity is presently making it feasible to store massive quantities of minimally processed machine data for later retrieval and analysis. In general, storing minimally processed machine data and performing analysis operations at search time can provide greater flexibility because it enables an analyst to search all of the machine data, instead of searching only a pre-specified set of data items. This may enable an analyst to investigate different aspects of the machine data that previously were unavailable for analysis.

(22) However, analyzing and searching massive quantities of machine data presents a number of challenges. For example, a data center, servers, or network appliances may generate many different types and formats of machine data (e.g., system logs, network packet data (e.g., wire data, etc.), sensor data, application program data, error logs, stack traces, system performance data, operating system data, virtualization data, etc.) from thousands of different components, which can collectively be very time-consuming to analyze. In another example, mobile devices may generate large amounts of information relating to data accesses, application performance, operating system performance, network performance, etc. There can be millions of mobile devices that concurrently report these types of information.

(23) These challenges can be addressed by using an event-based data intake and query system, such as the SPLUNK® ENTERPRISE, SPLUNK® CLOUD, or SPLUNK® CLOUD SERVICE system developed by Splunk Inc. of San Francisco, California. These systems represent the leading platform for providing real-time operational intelligence that enables organizations to collect,

index, and search machine data from various websites, applications, servers, networks, and mobile devices that power their businesses. The data intake and query system is particularly useful for analyzing data which is commonly found in system log files, network data, metrics data, tracing data, and other data input sources.

(24) In the data intake and query system, machine data is collected and stored as “events.” An event comprises a portion of machine data and is associated with a specific point in time. The portion of machine data may reflect activity in an IT environment and may be produced by a component of that IT environment, where the events may be searched to provide insight into the IT environment, thereby improving the performance of components in the IT environment. Events may be derived from “time series data,” where the time series data comprises a sequence of data points (e.g., performance measurements from a computer system, etc.) that are associated with successive points in time. In general, each event has a portion of machine data that is associated with a timestamp. The time stamp may be derived from the portion of machine data in the event, determined through interpolation between temporally proximate events having known timestamps, and/or may be determined based on other configurable rules for associating timestamps with events.

(25) In some instances, machine data can have a predefined structure, where data items with specific data formats are stored at predefined locations in the data. For example, the machine data may include data associated with fields in a database table. In other instances, machine data may not have a predefined structure (e.g., may not be at fixed, predefined locations), but may have repeatable (e.g., non-random) patterns. This means that some machine data can comprise various data items of different data types that may be stored at different locations within the data. For example, when the data source is an operating system log, an event can include one or more lines from the operating system log containing machine data that includes different types of performance and diagnostic information associated with a specific point in time (e.g., a timestamp).

(26) Examples of components which may generate machine data from which events can be derived include, but are not limited to, web servers, application servers, databases, firewalls, routers, operating systems, and software applications that execute on computer systems, mobile devices, sensors, Internet of Things (IoT) devices, etc. The machine data generated by such data sources can include, for example and without limitation, server log files, activity log files, configuration files, messages, network packet data, performance measurements, sensor measurements, etc.

(27) The data intake and query system can use flexible schema to specify how to extract information from events. A flexible schema may be developed and redefined as needed. The flexible schema can be applied to events “on the fly,” when it is needed (e.g., at search time, index time, ingestion time, etc.). When the schema is not applied to events until search time, the schema may be referred to as a “late-binding schema.”

(28) During operation, the data intake and query system receives machine data from any type and number of sources (e.g., one or more system logs, streams of network packet data, sensor data, application program data, error logs, stack traces, system performance data, etc.). The system parses the machine data to produce events each having a portion of machine data associated with a timestamp, and stores the events. The system enables users to run queries against the stored events to, for example, retrieve events that meet filter criteria specified in a query, such as criteria indicating certain keywords or having specific values in defined fields. Additional query terms can further process the event data, such as, by transforming the data, etc.

(29) As used herein, the term “field” can refer to a location in the machine data of an event containing one or more values for a specific data item. A field may be referenced by a field name associated with the field. As will be described in more detail herein, in some cases, a field is defined by an extraction rule (e.g., a regular expression) that derives one or more values or a sub-portion of text from the portion of machine data in each event to produce a value for the field for that event. The set of values produced are semantically-related (such as IP address), even though

the machine data in each event may be in different formats (e.g., semantically-related values may be in different positions in the events derived from different sources).

(30) As described above, the system stores the events in a data store. The events stored in the data store are field-searchable, where field-searchable herein refers to the ability to search the machine data (e.g., the raw machine data) of an event based on a field specified in search criteria. For example, a search having criteria that specifies a field name “UserID” may cause the system to field-search the machine data of events to identify events that have the field name “UserID.” In another example, a search having criteria that specifies a field name “UserID” with a corresponding field value “12345” may cause the system to field-search the machine data of events to identify events having that field-value pair (e.g., field name “UserID” with a corresponding field value of “12345”). Events are field-searchable using one or more configuration files associated with the events. Each configuration file can include one or more field names, where each field name is associated with a corresponding extraction rule and a set of events to which that extraction rule applies. The set of events to which an extraction rule applies may be identified by metadata associated with the set of events. For example, an extraction rule may apply to a set of events that are each associated with a particular host, source, or sourcetype. When events are to be searched based on a particular field name specified in a search, the system can use one or more configuration files to determine whether there is an extraction rule for that particular field name that applies to each event that falls within the criteria of the search. If so, the event is considered as part of the search results (and additional processing may be performed on that event based on criteria specified in the search). If not, the next event is similarly analyzed, and so on.

(31) As noted above, the data intake and query system can utilize a late-binding schema while performing queries on events. One aspect of a late-binding schema is applying extraction rules to events to extract values for specific fields during search time. More specifically, the extraction rule for a field can include one or more instructions that specify how to extract a value for the field from an event. An extraction rule can generally include any type of instruction for extracting values from machine data or events. In some cases, an extraction rule comprises a regular expression, where a sequence of characters form a search pattern. An extraction rule comprising a regular expression is referred to herein as a regex rule. The system applies a regex rule to machine data or an event to extract values for a field associated with the regex rule, where the values are extracted by searching the machine data/event for the sequence of characters defined in the regex rule.

(32) In the data intake and query system, a field extractor may be configured to automatically generate extraction rules for certain fields in the events when the events are being created, indexed, or stored, or possibly at a later time. Alternatively, a user may manually define extraction rules for fields using a variety of techniques. In contrast to a conventional schema for a database system, a late-binding schema is not defined at data ingestion time. Instead, the late-binding schema can be developed on an ongoing basis until the time a query is actually executed. This means that extraction rules for the fields specified in a query may be provided in the query itself, or may be located during execution of the query. Hence, as a user learns more about the data in the events, the user can continue to refine the late-binding schema by adding new fields, deleting fields, or modifying the field extraction rules for use the next time the schema is used by the system. Because the data intake and query system maintains the underlying machine data and uses a late-binding schema for searching the machine data, it enables a user to continue investigating and learn valuable insights about the machine data.

(33) In some embodiments, a common field name may be used to reference two or more fields containing equivalent and/or similar data items, even though the fields may be associated with different types of events that possibly have different data formats and different extraction rules. By enabling a common field name to be used to identify equivalent and/or similar fields from different types of events generated by disparate data sources, the system facilitates use of a “common information model” (CIM) across the disparate data sources.

(34) In some embodiments, the configuration files and/or extraction rules described above can be stored in a catalog, such as a metadata catalog. In certain embodiments, the content of the extraction rules can be stored as rules or actions in the metadata catalog. For example, the identification of the data to which the extraction rule applies can be referred to a rule and the processing of the data can be referred to as an action.

(35) 1.0. Operating Environment

(36) FIG. 1 is a block diagram of an embodiment of a data processing environment **100**. In the illustrated embodiment, the environment **100** includes a data intake and query system **102**, one or more host devices **104**, and one or more client computing devices **106** (generically referred to as client device(s) **106**).

(37) The data intake and query system **102**, host devices **104**, and client devices **106** can communicate with each other via one or more networks, such as a local area network (LAN), wide area network (WAN), private or personal network, cellular networks, intranetworks, and/or internetworks using any of wired, wireless, terrestrial microwave, satellite links, etc., and may include the Internet. Although not explicitly shown in FIG. 1, it will be understood that a client computing device **106** can communicate with a host device **104** via one or more networks. For example, if the host device **104** is configured as a web server and the client computing device **106** is a laptop, the laptop can communicate with the web server to view a website.

(38) 1.1. Client Devices

(39) A client device **106** can correspond to a distinct computing device that can configure, manage, or sends queries to the system **102**. Examples of client devices **106** may include, without limitation, smart phones, tablet computers, handheld computers, wearable devices, laptop computers, desktop computers, servers, portable media players, gaming devices, or other device that includes computer hardware (e.g., processors, non-transitory, computer-readable media, etc.) and so forth. In certain cases, a client device **106** can include a hosted, virtualized, or containerized device, such as an isolated execution environment, that shares computing resources (e.g., processor, memory, etc.) of a particular machine with other isolated execution environments.

(40) The client devices **106** can interact with the system **102** (or a host device **104**) in a variety of ways. For example, the client devices **106** can communicate with the system **102** (or a host device **104**) over an Internet (Web) protocol, via a gateway, via a command line interface, via a software developer kit (SDK), a standalone application, etc. As another example, the client devices **106** can use one or more executable applications or programs to interface with the system **102**.

(41) 1.2. Host Devices

(42) A host device **104** can correspond to a distinct computing device or system that includes or has access to data that can be ingested, indexed, and/or searched by the system **102**. Accordingly, in some cases, a client device **106** may also be a host device **104** (e.g., it can include data that is ingested by the system **102** and it can submit queries to the system **102**). The host devices **104** can include, but are not limited to, servers, sensors, routers, personal computers, mobile devices, internet of things (IoT) devices, or hosting devices, such as computing devices in a shared computing resource environment on which multiple isolated execution environment (e.g., virtual machines, containers, etc.) can be instantiated, or other computing devices in an IT environment (e.g., device that includes computer hardware, e.g., processors, non-transitory, computer-readable media, etc.). In certain cases, a host device **104** can include a hosted, virtualized, or containerized device, such as an isolated execution environment, that shares computing resources (e.g., processor, memory, etc.) of a particular machine (e.g., a hosting device or hosting machine) with other isolated execution environments.

(43) As mentioned host devices **104** can include or have access to data sources for the system **102**. The data sources can include machine data found in log files, data files, distributed file systems, streaming data, publication-subscribe (pub/sub) buffers, directories of files, data sent over a network, event logs, registries, streaming data services (examples of which can include, by way of

non-limiting example, AMAZON SIMPLE QUEUE SERVICE® (“SQS”) or Kinesis™ services, devices executing Apache Kafka™ software, or devices implementing the Message Queue Telemetry Transport (MQTT) protocol, MICROSOFT AZURE EVENTHUB®, Google Cloud™ Pub/Sub, devices implementing the Java Message Service (JMS) protocol, devices implementing the Advanced Message Queuing Protocol (AMQP)), cloud-based services (e.g., AWS®, MICROSOFT AZURE®, Google Cloud™, etc.), operating-system-level virtualization environments (e.g., Docker), container orchestration systems (e.g., Kubernetes), virtual machines using full virtualization or paravirtualization, or other virtualization technique or isolated execution environments.

(44) In some cases, one or more applications executing on a host device may generate various types of machine data during operation. For example, a web server application executing on a host device **104** may generate one or more web server logs detailing interactions between the web server and any number of client devices **106** or other devices. As another example, a host device **104** implemented as a router may generate one or more router logs that record information related to network traffic managed by the router. As yet another example, a database server application executing on a host device **104** may generate one or more logs that record information related to requests sent from other devices (e.g., web servers, application servers, client devices, etc.) for data managed by the database server. Similarly, a host device **104** may generate and/or store computing resource utilization metrics, such as, but not limited to, CPU utilization, memory utilization, number of processes being executed, etc. Any one or any combination of the files or data generated in such cases can be used as a data source for the system **102**.

(45) In some embodiments, an application may include a monitoring component that facilitates generating performance data related to host device's operating state, including monitoring network traffic sent and received from the host device and collecting other device and/or application-specific information. A monitoring component may be an integrated component of the application, a plug-in, an extension, or any other type of add-on component, or a stand-alone process.

(46) Such monitored information may include, but is not limited to, network performance data (e.g., a URL requested, a connection type (e.g., HTTP, HTTPS, etc.), a connection start time, a connection end time, an HTTP status code, request length, response length, request headers, response headers, connection status (e.g., completion, response time(s), failure, etc.)) or device performance information (e.g., current wireless signal strength of the device, a current connection type and network carrier, current memory performance information, processor utilization, memory utilization, a geographic location of the device, a device orientation, and any other information related to the operational state of the host device, etc.), device profile information (e.g., a type of client device, a manufacturer, and model of the device, versions of various software applications installed on the device, etc.) In some cases, the monitoring component can collect device performance information by monitoring one or more host device operations, or by making calls to an operating system and/or one or more other applications executing on a host device for performance information. The monitored information may be stored in one or more files and/or streamed to the system **102**.

(47) In general, a monitoring component may be configured to generate performance data in response to a monitor trigger in the code of a client application or other triggering application event, as described above, and to store the performance data in one or more data records. Each data record, for example, may include a collection of field-value pairs, each field-value pair storing a particular item of performance data in association with a field for the item. For example, a data record generated by a monitoring component may include a “networkLatency” field (not shown in the Figure) in which a value is stored. This field indicates a network latency measurement associated with one or more network requests. The data record may include a “state” field to store a value indicating a state of a network connection, and so forth for any number of aspects of collected performance data. Examples of functionality that enables monitoring performance of a

host device are described in U.S. patent application Ser. No. 14/524,748, entitled “UTILIZING PACKET HEADERS TO MONITOR NETWORK TRAFFIC IN ASSOCIATION WITH A CLIENT DEVICE,” filed on 27 Oct. 2014, and which is hereby incorporated by reference in its entirety for all purposes.

(48) In some embodiments, such as in a shared computing resource environment (or hosted environment), a host device **104** may include logs or machine data generated by an application executing within an isolated execution environment (e.g., web server log file if the isolated execution environment is configured as a web server or database server log files if the isolated execution environment is configured as database server, etc.), machine data associated with the computing resources assigned to the isolated execution environment (e.g., CPU utilization of the portion of the CPU allocated to the isolated execution environment, memory utilization of the portion of the memory allocated to the isolated execution environment, etc.), logs or machine data generated by an application that enables the isolated execution environment to share resources with other isolated execution environments (e.g., logs generated by a Docker manager or Kubernetes manager executing on the host device **104**), and/or machine data generated by monitoring the computing resources of the host device **104** (e.g., CPU utilization, memory utilization, etc.) that are shared between the isolated execution environments. Given the separation (and isolation) between isolated execution environments executing on a common computing device, in certain embodiments, each isolated execution environment may be treated as a separate host device **104** even if they are, in fact, executing on the same computing device or hosting device.

(49) Accordingly, as used herein, obtaining data from a data source may refer to communicating with a host device **104** to obtain data from the host device **104** (e.g., from one or more data source files, data streams, directories on the host device **104**, etc.). For example, obtaining data from a data source may refer to requesting data from a host device **104** and/or receiving data from a host device **104**. In some such cases, the host device **104** can retrieve and return the requested data from a particular data source and/or the system **102** can retrieve the data from a particular data source of the host device **104** (e.g., from a particular file stored on a host device **104**).

(50) 1.4. Data Intake and Query System Overview

(51) The data intake and query system **102** can ingest, index, and/or store data from heterogeneous data sources and/or host devices **104**. For example, the system **102** can ingest, index, and/or store any type of machine data, regardless of the form of the machine data or whether the machine data matches or is similar to other machine data ingested, indexed, and/or stored by the system **102**. In some cases, the system **102** can generate events from the received data, group the events, and store the events in buckets. The system **102** can also search heterogeneous data that it has stored or search data stored by other systems (e.g., other system **102** systems or other non-system **102** systems). For example, in response to received queries, the system **102** can assign one or more components to search events stored in the storage system or search data stored elsewhere.

(52) As will be described herein in greater detail below, the system **102** can use one or more components to ingest, index, store, and/or search data. In some embodiments, the system **102** is implemented as a distributed system that uses multiple components to perform its various functions. For example, the system **102** can include any one or any combination of subsystems (referenced as “systems”) such as an intake system **110** (including one or more components) to ingest data, an indexing system **112** (including one or more components) to index the data, a storage system **116** (including one or more components) to store the data, and/or a query system **114** (including one or more components) to search the data, etc.

(53) In the illustrated embodiment, the system **102** is shown having four subsystems **110**, **112**, **114**, **116**. However, it will be understood that the system **102** may include any one or any combination of the intake system **110**, indexing system **112**, query system **114**, or storage system **116**. Further, in certain embodiments, one or more of the intake system **110**, indexing system **112**, query system **114**, or storage system **116** may be used alone or apart from the system **102**. For example, the

intake system **110** may be used alone to glean information from streaming data that is not indexed or stored by the system **102**, or the query system **114** may be used to search data that is unaffiliated with the system **102**.

(54) In certain embodiments, the components of the different systems may be distinct from each other or there may be some overlap. For example, one component of the system **102** may include some indexing functionality and some searching functionality and thus be used as part of the indexing system **112** and query system **114**, while another computing device of the system **102** may only have ingesting or search functionality and only be used as part of those respective systems. Similarly, the components of the storage system **116** may include data stores of individual components of the indexing system and/or may be a separate shared data storage system, like AMAZON SIMPLE STORAGE SERVICE® (S3), that is accessible to distinct components of the intake system **110**, indexing system **112**, and query system **114**.

(55) In some cases, the components of the system **102** are implemented as distinct computing devices having their own computer hardware (e.g., processors, non-transitory, computer-readable media, etc.) and/or as distinct hosted devices (e.g., isolated execution environments) that share computing resources or hardware in a shared computing resource environment. However, the components of the system **102**, such as the query system **114** for example, may be implemented as logical devices (e.g., instances, etc.). These logical devices perform certain functionality or functionalities of a counterpart physical (hardware-based) device based on processor execution.

(56) For simplicity, references made herein to the intake system **110**, indexing system **112**, storage system **116**, and query system **114** can refer to those components used for ingesting, indexing, storing, and searching, respectively. However, it will be understood that although reference is made to two separate systems, the same underlying component may be performing the functions for the two different systems. For example, reference to the indexing system indexing data and storing the data in the storage system **116** or the query system searching the data may refer to the same component (e.g., same computing device or hosted device) indexing the data, storing the data, and then searching the data that it stored.

(57) 1.4.1. Intake System Overview

(58) As will be described in greater detail herein, the intake system **110** can receive data from the host devices **104** or data sources, perform one or more preliminary processing operations on the data, and communicate the data to the indexing system **112**, query system **114**, storage system **116**, or to other systems (which may include, for example, data processing systems, telemetry systems, real-time analytics systems, data stores, databases, etc., any of which may be operated by an operator of the system **102** or a third party). Given the amount of data that can be ingested by the intake system **110**, in some embodiments, the intake system can include multiple distributed computing devices or components working concurrently to ingest the data.

(59) The intake system **110** can receive data from the host devices **104** in a variety of formats or structures. In some embodiments, the received data corresponds to raw machine data, structured or unstructured data, correlation data, data files, directories of files, data sent over a network, event logs, registries, messages published to streaming data sources, performance metrics, sensor data, image and video data, etc.

(60) The preliminary processing operations performed by the intake system **110** can include, but is not limited to, associating metadata with the data received from a host device **104**, extracting a timestamp from the data, identifying individual events within the data, extracting a subset of machine data for transmittal to the indexing system **112**, enriching the data, etc. As part of communicating the data to the indexing system, the intake system **110** can route the data to a particular component of the intake system **110** or dynamically route the data based on load-balancing, etc. In certain cases, one or more components of the intake system **110** can be installed on a host device **104**.

(61) 1.4.2. Indexing System Overview

(62) As will be described in greater detail herein, the indexing system **112** can include one or more components (e.g., indexing nodes) to process the data and store it, for example, in the storage system **116**. As part of processing the data, the indexing system can identify distinct events within the data, timestamps associated with the data, organize the data into buckets or time series buckets, convert editable buckets to non-editable buckets, store copies of the buckets in the storage system **116**, merge buckets, generate indexes of the data, etc. In addition, the indexing system **112** can update various catalogs or databases with information related to the buckets (pre-merged or merged) or data that is stored in the storage system **116**, and can communicate with the intake system **110** about the status of the data storage.

(63) 1.4.3. Query System Overview

(64) As will be described in greater detail herein, the query system **114** can include one or more components to receive, process, and execute queries. In some cases, the query system **114** can use the same component to process and execute the query or use one or more components to receive and process the query (e.g., a search head) and use one or more other components to execute at least a portion of the query (e.g., search nodes). In some cases, a search node and an indexing node may refer to the same computing device or hosted device performing different functions. In certain cases, a search node can be a separate computing device or hosted device from an indexing node.

(65) Queries received by the query system **114** can be relatively complex and identify a set of data to be processed and a manner of processing the set of data from one or more client devices **106**. In certain cases, the query can be implemented using a pipelined command language or other query language. As described herein, in some cases, the query system **114** can execute parts of the query in a distributed fashion (e.g., one or more mapping phases or parts associated with identifying and gathering the set of data identified in the query) and execute other parts of the query on a single component (e.g., one or more reduction phases). However, it will be understood that in some cases multiple components can be used in the map and/or reduce functions of the query execution.

(66) In some cases, as part of executing the query, the query system **114** can use one or more catalogs or databases to identify the set of data to be processed or its location in the storage system **116** and/or can retrieve data from the storage system **116**. In addition, in some embodiments, the query system **114** can store some or all of the query results in the storage system **116**.

(67) 1.4.4. Storage System Overview

(68) In some cases, the storage system **116** may include one or more data stores associated with or coupled to the components of the indexing system **112** that are accessible via a system bus or local area network. In certain embodiments, the storage system **116** may be a shared storage system **116**, like AMAZON S3® or Google Cloud Storage™, that are accessible via a wide area network.

(69) As mentioned and as will be described in greater detail below, the storage system **116** can be made up of one or more data stores storing data that has been processed by the indexing system **112**. In some cases, the storage system includes data stores of the components of the indexing system **112** and/or query system **114**. In certain embodiments, the storage system **116** can be implemented as a shared storage system **116**. The shared storage system **116** can be configured to provide high availability, highly resilient, low loss data storage. In some cases, to provide the high availability, highly resilient, low loss data storage, the shared storage system **116** can store multiple copies of the data in the same and different geographic locations and across different types of data stores (e.g., solid state, hard drive, tape, etc.). Further, as data is received at the shared storage system **116** it can be automatically replicated multiple times according to a replication factor to different data stores across the same and/or different geographic locations. In some embodiments, the shared storage system **116** can correspond to cloud storage, such as Amazon S3® or Elastic Block Storage (EBS), Google Cloud™ Storage, MICROSOFT AZURE® Storage, etc.

(70) In some embodiments, indexing system **112** can read to and write from the shared storage system **116**. For example, the indexing system **112** can copy buckets of data from its local or shared data stores to the shared storage system **116**. In certain embodiments, the query system **114** can

read from, but cannot write to, the shared storage system **116**. For example, the query system **114** can read the buckets of data stored in shared storage system **116** by the indexing system **112**, but may not be able to copy buckets or other data to the shared storage system **116**. In some embodiments, the intake system **110** does not have access to the shared storage system **116**. However, in some embodiments, one or more components of the intake system **110** can write data to the shared storage system **116** that can be read by the indexing system **112**.

(71) As described herein, in some embodiments, data in the system **102** (e.g., in the data stores of the components of the indexing system **112**, shared storage system **116**, or search nodes of the query system **114**) can be stored in one or more time series buckets. Each bucket can include raw machine data associated with a timestamp and additional information about the data or bucket, such as, but not limited to, one or more filters, indexes (e.g., TSIDX, inverted indexes, keyword indexes, etc.), bucket summaries, etc. In some embodiments, the bucket data and information about the bucket data is stored in one or more files. For example, the raw machine data, filters, indexes, bucket summaries, etc. can be stored in respective files in or associated with a bucket. In certain cases, the group of files can be associated together to form the bucket.

(72) 1.4.5. Other Components of the Data Intake and Query System

(73) The system **102** can include additional components that interact with any one or any combination of the intake system **110**, indexing system **112**, query system **114**, and/or storage system **116**. Such components may include, but are not limited to an authentication system, orchestration system, one or more catalogs or databases, a gateway, etc.

(74) An authentication system can include one or more components to authenticate users to access, use, and/or configure the system **102**. Similarly, the authentication system can be used to restrict what a particular user can do on the system **102** and/or what components or data a user can access, etc.

(75) An orchestration system can include one or more components to manage and/or monitor the various components of the system **102**. In some embodiments, the orchestration system can monitor the components of the system **102** to detect when one or more components has failed or is unavailable and enable the system **102** to recover from the failure (e.g., by adding additional components, fixing the failed component, or having other components complete the tasks assigned to the failed component). In certain cases, the orchestration system can determine when to add components to or remove components from a particular system **110**, **112**, **114**, **116** (e.g., based on usage, user/tenant requests, etc.). In embodiments where the system **102** is implemented in a shared computing resource environment, the orchestration system can facilitate the creation and/or destruction of isolated execution environments or instances of the components of the system **102**, etc.

(76) In certain embodiments, the system **102** can include various components that enable it to provide stateless services or enable it to recover from an unavailable or unresponsive component without data loss in a time efficient manner. For example, the system **102** can store contextual information about its various components in a distributed way such that if one of the components becomes unresponsive or unavailable, the system **102** can replace the unavailable component with a different component and provide the replacement component with the contextual information. In this way, the system **102** can quickly recover from an unresponsive or unavailable component while reducing or eliminating the loss of data that was being processed by the unavailable component.

(77) In some embodiments, the system **102** can store the contextual information in a catalog, as described herein. In certain embodiments, the contextual information can correspond to information that the system **102** has determined or learned based on use. In some cases, the contextual information can be stored as annotations (manual annotations and/or system annotations), as described herein.

(78) In certain embodiments, the system **102** can include an additional catalog that monitors the

location and storage of data in the storage system **116** to facilitate efficient access of the data during search time. In certain embodiments, such a catalog may form part of the storage system **116**.

(79) In some embodiments, the system **102** can include a gateway or other mechanism to interact with external devices or to facilitate communications between components of the system **102**. In some embodiments, the gateway can be implemented using an application programming interface (API). In certain embodiments, the gateway can be implemented using a representational state transfer API (REST API).

(80) 1.4.6. On-Premise and Shared Computing Resource Environments

(81) In some environments, a user of a system **102** may install and configure, on computing devices owned and operated by the user, one or more software applications that implement some or all of the components of the system **102**. For example, with reference to FIG. **1**, a user may install a software application on server computers owned by the user and configure each server to operate as one or more components of the intake system **110**, indexing system **112**, query system **114**, shared storage system **116**, or other components of the system **102**. This arrangement generally may be referred to as an “on-premises” solution. That is, the system **102** is installed and operates on computing devices directly controlled by the user of the system **102**. Some users may prefer an on-premises solution because it may provide a greater level of control over the configuration of certain aspects of the system (e.g., security, privacy, standards, controls, etc.). However, other users may instead prefer an arrangement in which the user is not directly responsible for providing and managing the computing devices upon which various components of system **102** operate.

(82) In certain embodiments, one or more of the components of the system **102** can be implemented in a shared computing resource environment. In this context, a shared computing resource environment or cloud-based service can refer to a service hosted by one more computing resources that are accessible to end users over a network, for example, by using a web browser or other application on a client device to interface with the remote computing resources. For example, a service provider may provide a system **102** by managing computing resources configured to implement various aspects of the system (e.g., intake system **110**, indexing system **112**, query system **114**, shared storage system **116**, other components, etc.) and by providing access to the system to end users via a network. Typically, a user may pay a subscription or other fee to use such a service. Each subscribing user of the cloud-based service may be provided with an account that enables the user to configure a customized cloud-based system based on the user's preferences.

(83) When implemented in a shared computing resource environment, the underlying hardware (non-limiting examples: processors, hard drives, solid-state memory, RAM, etc.) on which the components of the system **102** execute can be shared by multiple customers or tenants as part of the shared computing resource environment. In addition, when implemented in a shared computing resource environment as a cloud-based service, various components of the system **102** can be implemented using containerization or operating-system-level virtualization, or other virtualization technique. For example, one or more components of the intake system **110**, indexing system **112**, or query system **114** can be implemented as separate software containers or container instances. Each container instance can have certain computing resources (e.g., memory, processor, etc.) of an underlying hosting computing system (e.g., server, microprocessor, etc.) assigned to it, but may share the same operating system and may use the operating system's system call interface. Each container may provide an isolated execution environment on the host system, such as by providing a memory space of the hosting system that is logically isolated from memory space of other containers. Further, each container may run the same or different computer applications concurrently or separately, and may interact with each other. Although reference is made herein to containerization and container instances, it will be understood that other virtualization techniques can be used. For example, the components can be implemented using virtual machines using full virtualization or paravirtualization, etc. Thus, where reference is made to “containerized” components, it should be understood that such components may additionally or alternatively be

implemented in other isolated execution environments, such as a virtual machine environment.

(84) Implementing the system **102** in a shared computing resource environment can provide a number of benefits. In some cases, implementing the system **102** in a shared computing resource environment can make it easier to install, maintain, and update the components of the system **102**. For example, rather than accessing designated hardware at a particular location to install or provide a component of the system **102**, a component can be remotely instantiated or updated as desired. Similarly, implementing the system **102** in a shared computing resource environment or as a cloud-based service can make it easier to meet dynamic demand. For example, if the system **102** experiences significant load at indexing or search, additional compute resources can be deployed to process the additional data or queries. In an “on-premises” environment, this type of flexibility and scalability may not be possible or feasible.

(85) In addition, by implementing the system **102** in a shared computing resource environment or as a cloud-based service can improve compute resource utilization. For example, in an on-premises environment if the designated compute resources are not being used by, they may sit idle and unused. In a shared computing resource environment, if the compute resources for a particular component are not being used, they can be re-allocated to other tasks within the system **102** and/or to other systems unrelated to the system **102**.

(86) As mentioned, in an on-premises environment, data from one instance of a system **102** is logically and physically separated from the data of another instance of a system **102** by virtue of each instance having its own designated hardware. As such, data from different customers of the system **102** is logically and physically separated from each other. In a shared computing resource environment, components of a system **102** can be configured to process the data from one customer or tenant or from multiple customers or tenants. Even in cases where a separate component of a system **102** is used for each customer, the underlying hardware on which the components of the system **102** are instantiated may still process data from different tenants. Accordingly, in a shared computing resource environment, the data from different tenants may not be physically separated on distinct hardware devices. For example, data from one tenant may reside on the same hard drive as data from another tenant or be processed by the same processor. In such cases, the system **102** can maintain logical separation between tenant data. For example, the system **102** can include separate directories for different tenants and apply different permissions and access controls to access the different directories or to process the data, etc.

(87) In certain cases, the tenant data from different tenants is mutually exclusive and/or independent from each other. For example, in certain cases, Tenant A and Tenant B do not share the same data, similar to the way in which data from a local hard drive of Customer A is mutually exclusive and independent of the data (and not considered part) of a local hard drive of Customer B. While Tenant A and Tenant B may have matching or identical data, each tenant would have a separate copy of the data. For example, with reference again to the local hard drive of Customer A and Customer B example, each hard drive could include the same file. However, each instance of the file would be considered part of the separate hard drive and would be independent of the other file. Thus, one copy of the file would be part of Customer's A hard drive and a separate copy of the file would be part of Customer B's hard drive. In a similar manner, to the extent Tenant A has a file that is identical to a file of Tenant B, each tenant would have a distinct and independent copy of the file stored in different locations on a data store or on different data stores.

(88) Further, in certain cases, the system **102** can maintain the mutual exclusivity and/or independence between tenant data even as the tenant data is being processed, stored, and searched by the same underlying hardware. In certain cases, to maintain the mutual exclusivity and/or independence between the data of different tenants, the system **102** can use tenant identifiers to uniquely identify data associated with different tenants.

(89) In a shared computing resource environment, some components of the system **102** can be instantiated and designated for individual tenants and other components can be shared by multiple

tenants. In certain embodiments, a separate intake system **110**, indexing system **112**, and query system **114** can be instantiated for each tenant, whereas the shared storage system **116** or other components (e.g., data store, metadata catalog, and/or acceleration data store, described below) can be shared by multiple tenants. In some such embodiments where components are shared by multiple tenants, the components can maintain separate directories for the different tenants to ensure their mutual exclusivity and/or independence from each other. Similarly, in some such embodiments, the system **102** can use different hosting computing systems or different isolated execution environments to process the data from the different tenants as part of the intake system **110**, indexing system **112**, and/or query system **114**.

(90) In some embodiments, individual components of the intake system **110**, indexing system **112**, and/or query system **114** may be instantiated for each tenant or shared by multiple tenants. For example, some individual intake system components (e.g., forwarders, output ingestion buffer) may be instantiated and designated for individual tenants, while other intake system components (e.g., a data retrieval subsystem, intake ingestion buffer, and/or streaming data processor), may be shared by multiple tenants.

(91) In certain embodiments, an indexing system **112** (or certain components thereof) can be instantiated and designated for a particular tenant or shared by multiple tenants. In some embodiments where a separate indexing system **112** is instantiated and designated for each tenant, different resources can be reserved for different tenants. For example, Tenant A can be consistently allocated a minimum of four indexing nodes and Tenant B can be consistently allocated a minimum of two indexing nodes. In some such embodiments, the four indexing nodes can be reserved for Tenant A and the two indexing nodes can be reserved for Tenant B, even if Tenant A and Tenant B are not using the reserved indexing nodes.

(92) In embodiments where an indexing system **112** is shared by multiple tenants, components of the indexing system **112** can be dynamically assigned to different tenants. For example, if Tenant A has greater indexing demands, additional indexing nodes can be instantiated or assigned to Tenant A's data. However, as the demand decreases, the indexing nodes can be reassigned to a different tenant, or terminated. Further, in some embodiments, a component of the indexing system **112** can concurrently process data from the different tenants.

(93) In some embodiments, one instance of query system **114** may be shared by multiple tenants. In some such cases, the same search head can be used to process/execute queries for different tenants and/or the same search nodes can be used to execute query for different tenants. Further, in some such cases, different tenants can be allocated different amounts of compute resources. For example, Tenant A may be assigned more search heads or search nodes based on demand or based on a service level arrangement than another tenant. However, once a search is completed the search head and/or nodes assigned to Tenant A may be assigned to Tenant B, deactivated, or their resource may be re-allocated to other components of the system **102**, etc.

(94) In some cases, by sharing more components with different tenants, the functioning of the system **102** can be improved. For example, by sharing components across tenants, the system **102** can improve resource utilization thereby reducing the amount of resources allocated as a whole. For example, if four indexing nodes, two search heads, and four search nodes are reserved for each tenant then those compute resources are unavailable for use by other processes or tenants, even if they go unused. In contrast, by sharing the indexing nodes, search heads, and search nodes with different tenants and instantiating additional compute resources, the system **102** can use fewer resources overall while providing improved processing time for the tenants that are using the compute resources. For example, if tenant A is not using any search nodes and tenant B has many searches running, the system **102** can use search nodes that would have been reserved for tenant A to service tenant B. In this way, the system **102** can decrease the number of compute resources used/reserved, while improving the search time for tenant B and improving compute resource utilization.

(95) 2.0. Data Ingestion, Indexing, and Storage

(96) FIG. 2 is a flow diagram illustrating an embodiment of a routine implemented by the system **102** to process, index, and store data received from host devices **104**. The data flow illustrated in FIG. 2 is provided for illustrative purposes only. It will be understood that one or more of the steps of the processes illustrated in FIG. 2 may be removed or that the ordering of the steps may be changed. Furthermore, for the purposes of illustrating a clear example, one or more particular system components are described in the context of performing various operations during each of the data flow stages. For example, the intake system **110** is described as receiving machine data and the indexing system **112** is described as generating events, grouping events, and storing events. However, other system arrangements and distributions of the processing steps across system components may be used. For example, in some cases, the intake system **110** may generate events. (97) At block **202**, the intake system **110** receives data from a host device **104**. The intake system **110** initially may receive the data as a raw data stream generated by the host device **104**. For example, the intake system **110** may receive a data stream from a log file generated by an application server, from a stream of network data from a network device, or from any other source of data. Non-limiting examples of machine data that can be received by the intake system **110** is described herein with reference to FIG. 3A.

(98) In some embodiments, the intake system **110** receives the raw data and may segment the data stream into messages, possibly of a uniform data size, to facilitate subsequent processing steps. The intake system **110** may thereafter process the messages in accordance with one or more rules to conduct preliminary processing of the data. In one embodiment, the processing conducted by the intake system **110** may be used to indicate one or more metadata fields applicable to each message. For example, the intake system **110** may include metadata fields within the messages, or publish the messages to topics indicative of a metadata field. These metadata fields may, for example, provide information related to a message as a whole and may apply to each event that is subsequently derived from the data in the message. For example, the metadata fields may include separate fields specifying each of a host, a source, and a sourcetype related to the message. A host field may contain a value identifying a host name or IP address of a device that generated the data. A source field may contain a value identifying a source of the data, such as a pathname of a file or a protocol and port related to received network data. A sourcetype field may contain a value specifying a particular sourcetype label for the data. Additional metadata fields may also be included, such as a character encoding of the data, if known, and possibly other values that provide information relevant to later processing steps. In certain embodiments, the intake system **110** may perform additional operations, such as, but not limited to, identifying individual events within the data, determining timestamps for the data, further enriching the data, etc.

(99) At block **204**, the indexing system **112** generates events from the data. In some cases, as part of generating the events, the indexing system **112** can parse the data of the message. In some embodiments, the indexing system **112** can determine a sourcetype associated with each message (e.g., by extracting a sourcetype label from the metadata fields associated with the message, etc.) and refer to a sourcetype configuration corresponding to the identified sourcetype to parse the data of the message. The sourcetype definition may include one or more properties that indicate to the indexing system **112** to automatically determine the boundaries within the received data that indicate the portions of machine data for events. In general, these properties may include regular expression-based rules or delimiter rules where, for example, event boundaries may be indicated by predefined characters or character strings. These predefined characters may include punctuation marks or other special characters including, for example, carriage returns, tabs, spaces, line breaks, etc. If a sourcetype for the data is unknown to the indexing system **112**, the indexing system **112** may infer a sourcetype for the data by examining the structure of the data. Then, the indexing system **112** can apply an inferred sourcetype definition to the data to create the events.

(100) In addition, as part of generating events from the data, the indexing system **112** can determine

a timestamp for each event. Similar to the process for parsing machine data, the indexing system **112** may again refer to a sourcetype definition associated with the data to locate one or more properties that indicate instructions for determining a timestamp for each event. The properties may, for example, instruct the indexing system **112** to extract a time value from a portion of data for the event (e.g., using a regex rule), to interpolate time values based on timestamps associated with temporally proximate events, to create a timestamp based on a time the portion of machine data was received or generated, to use the timestamp of a previous event, or use any other rules for determining timestamps, etc.

(101) The indexing system **112** can also associate events with one or more metadata fields. In some embodiments, a timestamp may be included in the metadata fields. These metadata fields may include any number of “default fields” that are associated with all events, and may also include one more custom fields as defined by a user. In certain embodiments, the default metadata fields associated with each event may include a host, source, and sourcetype field including or in addition to a field storing the timestamp.

(102) In certain embodiments, the indexing system **112** can also apply one or more transformations to event data that is to be included in an event. For example, such transformations can include removing a portion of the event data (e.g., a portion used to define event boundaries, extraneous characters from the event, other extraneous text, etc.), masking a portion of event data (e.g., masking a credit card number), removing redundant portions of event data, etc. The transformations applied to event data may, for example, be specified in one or more configuration files and referenced by one or more sourcetype definitions.

(103) At block **206**, the indexing system **112** can group events. In some embodiments, the indexing system **112** can group events based on time. For example, events generated within a particular time period or events that have a time stamp within a particular time period can be grouped together to form a bucket. A non-limiting example of a bucket is described herein with reference to FIG. **3B**.

(104) In certain embodiments, multiple components of the indexing system, such as an indexing node, can concurrently generate events and buckets. Furthermore, each indexing node that generates and groups events can concurrently generate multiple buckets. For example, multiple processors of an indexing node can concurrently process data, generate events, and generate buckets. Further, multiple indexing nodes can concurrently generate events and buckets. As such, ingested data can be processed in a highly distributed manner.

(105) In some embodiments, as part of grouping events together, the indexing system **112** can generate one or more inverted indexes for a particular group of events. A non-limiting example of an inverted index is described herein with reference to FIG. **3C**. In certain embodiments, the inverted indexes can include location information for events of a bucket. For example, the events of a bucket may be compressed into one or more files to reduce their size. The inverted index can include location information indicating the particular file and/or location within a particular file of a particular event.

(106) In certain embodiments, the inverted indexes may include keyword entries or entries for field values or field name-value pairs found in events. In some cases, a field name-value pair can include a pair of words connected by a symbol, such as an equals sign or colon. The entries can also include location information for events that include the keyword, field value, or field value pair. In this way, relevant events can be quickly located. In some embodiments, fields can automatically be generated for some or all of the field names of the field name-value pairs at the time of indexing. For example, if the string “dest=10.0.1.2” is found in an event, a field named “dest” may be created for the event, and assigned a value of “10.0.1.2.” In certain embodiments, the indexing system can populate entries in the inverted index with field name-value pairs by parsing events using one or more regex rules to determine a field value associated with a field defined by the regex rule. For example, the regex rule may indicate how to find a field value for a UserID field in certain events. In some cases, the indexing system **112** can use the sourcetype of the event to determine which

regex to use for identifying field values.

(107) At block **208**, the indexing system **112** stores the events with an associated timestamp in the storage system **116**, which may be in a local data store and/or in a shared storage system.

Timestamps enable a user to search for events based on a time range. In some embodiments, the stored events are organized into “buckets,” where each bucket stores events associated with a specific time range based on the timestamps associated with each event. As mentioned, FIGS. **3B** and **3C** illustrate an example of a bucket. This improves time-based searching, as well as allows for events with recent timestamps, which may have a higher likelihood of being accessed, to be stored in a faster memory to facilitate faster retrieval. For example, buckets containing the most recent events can be stored in flash memory rather than on a hard disk. In some embodiments, each bucket may be associated with an identifier, a time range, and a size constraint.

(108) The indexing system **112** may be responsible for storing the events in the storage system **116**. As mentioned, the events or buckets can be stored locally on a component of the indexing system **112** or in a shared storage system **116**. In certain embodiments, the component that generates the events and/or stores the events (indexing node) can also be assigned to search the events. In some embodiments separate components can be used for generating and storing events (indexing node) and for searching the events (search node).

(109) By storing events in a distributed manner (either by storing the events at different components or in a shared storage system **116**), the query system **114** can analyze events for a query in parallel. For example, using map-reduce techniques, multiple components of the query system (e.g., indexing or search nodes) can concurrently search and provide partial responses for a subset of events to another component (e.g., search head) that combines the results to produce an answer for the query. By storing events in buckets for specific time ranges, the indexing system **112** may further optimize the data retrieval process by the query system **114** to search buckets corresponding to time ranges that are relevant to a query. In some embodiments, each bucket may be associated with an identifier, a time range, and a size constraint. In certain embodiments, a bucket can correspond to a file system directory and the machine data, or events, of a bucket can be stored in one or more files of the file system directory. The file system directory can include additional files, such as one or more inverted indexes, high performance indexes, permissions files, configuration files, etc.

(110) In embodiments where components of the indexing system **112** store buckets locally, the components can include a home directory and a cold directory. The home directory can store hot buckets and warm buckets, and the cold directory stores cold buckets. A hot bucket can refer to a bucket that is capable of receiving and storing additional events. A warm bucket can refer to a bucket that can no longer receive events for storage, but has not yet been moved to the cold directory. A cold bucket can refer to a bucket that can no longer receive events and may be a bucket that was previously stored in the home directory. The home directory may be stored in faster memory, such as flash memory, as events may be actively written to the home directory, and the home directory may typically store events that are more frequently searched and thus are accessed more frequently. The cold directory may be stored in slower and/or larger memory, such as a hard disk, as events are no longer being written to the cold directory, and the cold directory may typically store events that are not as frequently searched and thus are accessed less frequently. In some embodiments, components of the indexing system **112** may also have a quarantine bucket that contains events having potentially inaccurate information, such as an incorrect timestamp associated with the event or a timestamp that appears to be an unreasonable timestamp for the corresponding event. The quarantine bucket may have events from any time range; as such, the quarantine bucket may always be searched at search time. Additionally, components of the indexing system may store old, archived data in a frozen bucket that is not capable of being searched at search time. In some embodiments, a frozen bucket may be stored in slower and/or larger memory, such as a hard disk, and may be stored in offline and/or remote storage.

(111) In some embodiments, components of the indexing system **112** may not include a cold directory and/or cold or frozen buckets. For example, in embodiments where buckets are copied to a shared storage system **116** and searched by separate components of the query system **114**, buckets can be deleted from components of the indexing system as they are stored to the storage system **116**. In certain embodiments, the shared storage system **116** may include a home directory that includes warm buckets copied from the indexing system **112** and a cold directory of cold or frozen buckets as described above.

(112) 2.1. Machine Data and Data Storage Examples

(113) FIG. **3A** is a block diagram illustrating an embodiment of machine data received by the system **102**. The machine data can correspond to data from one or more host devices **104** or data sources. As mentioned, the data source can correspond to a log file, data stream or other data structure that is accessible by a host device **104**. In the illustrated embodiment of FIG. **3A**, the machine data has different forms. For example, the machine data **302** may be log data that is unstructured or that does not have any clear structure or fields, and include different portions **302A-302E** that correspond to different entries of the log and that separated by boundaries. Such data may also be referred to as raw machine data.

(114) The machine data **304** may be referred to as structured or semi-structured machine data as it does include some data in a JSON structure defining certain field and field values (e.g., machine data **304A** showing field name: field values container_name: kube-apiserver, host: ip 172 20 43 173.ec2.internal, pod_id: 0a73017b-4efa-11e8-a4e1-0a2bf2ab4bba, etc.), but other parts of the machine data **304** is unstructured or raw machine data (e.g., machine data **304B**). The machine data **306** may be referred to as structured data as it includes particular rows and columns of data with field names and field values.

(115) In some embodiments, the machine data **302** can correspond to log data generated by a host device **104** configured as an Apache server, the machine data **304** can correspond to log data generated by a host device **104** in a shared computing resource environment, and the machine data **306** can correspond to metrics data. Given the differences between host devices **104** that generated the log data **302**, **304**, the form of the log data **302**, **304** is different. In addition, as the log data **304** is from a host device **104** in a shared computing resource environment, it can include log data generated by an application being executed within an isolated execution environment (**304B**, excluding the field name “log:”) and log data generated by an application that enables the sharing of computing resources between isolated execution environments (all other data in **304**). Although shown together in FIG. **3A**, it will be understood that machine data with different hosts, sources, or sourcetypes can be received separately and/or found in different data sources and/or host devices **104**.

(116) As described herein, the system **102** can process the machine data based on the form in which it is received. In some cases, the intake system **110** can utilize one or more rules to process the data. In certain embodiments, the intake system **110** can enrich the received data. For example, the intake system **110** may add one or more fields to the data received from the host devices **104**, such as fields denoting the host, source, sourcetype, index, or tenant associated with the incoming data. In certain embodiments, the intake system **110** can perform additional processing on the incoming data, such as transforming structured data into unstructured data (or vice versa), identifying timestamps associated with the data, removing extraneous data, parsing data, indexing data, separating data, categorizing data, routing data based on criteria relating to the data being routed, and/or performing other data transformations, etc.

(117) In some cases, the data processed by the intake system **110** can be communicated or made available to the indexing system **112**, the query system **114**, and/or to other systems. In some embodiments, the intake system **110** communicates or makes available streams of data using one or more shards. For example, the indexing system **112** may read or receive data from one shard and another system may receive data from another shard. As another example, multiple systems may

receive data from the same shard.

(118) As used herein, a partition can refer to a logical division of data. In some cases, the logical division of data may refer to a portion of a data stream, such as a shard from the intake system **110**. In certain cases, the logical division of data can refer to an index or other portion of data stored in the storage system **116**, such as different directories or file structures used to store data or buckets. Accordingly, it will be understood that the logical division of data referenced by the term partition will be understood based on the context of its use.

(119) FIGS. **3B** and **3C** are block diagrams illustrating embodiments of various data structures for storing data processed by the system **102**. FIG. **3B** includes an expanded view illustrating an example of machine data stored in a data store **310** of the data storage system **116**. It will be understood that the depiction of machine data and associated metadata as rows and columns in the table **319** of FIG. **3B** is merely illustrative and is not intended to limit the data format in which the machine data and metadata is stored in various embodiments described herein. In one particular embodiment, machine data can be stored in a compressed or encrypted format. In such embodiments, the machine data can be stored with or be associated with data that describes the compression or encryption scheme with which the machine data is stored. The information about the compression or encryption scheme can be used to decompress or decrypt the machine data, and any metadata with which it is stored, at search time.

(120) In the illustrated embodiment of FIG. **3B** the data store **310** includes a directory **312** (individually referred to as **312A**, **312B**) for each index (or partition) that contains a portion of data stored in the data store **310** and a sub-directory **314** (individually referred to as **314A**, **314B**, **314C**) for one or more buckets of the index. In the illustrated embodiment of FIG. **3B**, each sub-directory **314** corresponds to a bucket and includes an event data file **316** (individually referred to as **316A**, **316B**, **316C**) and an inverted index **318** (individually referred to as **318A**, **318B**, **318C**). However, it will be understood that each bucket can be associated with fewer or more files and each sub-directory **314** can store fewer or more files.

(121) In the illustrated embodiment, the data store **310** includes a main directory **312A** associated with an index “_main” and a _test directory **312B** associated with an index “test.” However, the data store **310** can include fewer or more directories. In some embodiments, multiple indexes can share a single directory or all indexes can share a common directory. Additionally, although illustrated as a single data store **310**, it will be understood that the data store **310** can be implemented as multiple data stores storing different portions of the information shown in FIG. **3C**. For example, a single index can span multiple directories or multiple data stores.

(122) Furthermore, although not illustrated in FIG. **3B**, it will be understood that, in some embodiments, the data store **310** can include directories for each tenant and sub-directories for each index of each tenant, or vice versa. Accordingly, the directories **312A** and **312B** can, in certain embodiments, correspond to sub-directories of a tenant or include sub-directories for different tenants.

(123) In the illustrated embodiment of FIG. **3B**, two sub-directories **314A**, **314B** of the main directory **312A** and one sub-directory **312C** of the _test directory **312B** are shown. The sub-directories **314A**, **314B**, **314C** can correspond to buckets of the indexes associated with the directories **312A**, **312B**. For example, the sub-directories **314A** and **314B** can correspond to buckets “B1” and “B2,” respectively, of the index “_main” and the sub-directory **314C** can correspond to bucket “B1” of the index “_test.” Accordingly, even though there are two “B1” buckets shown, as each “B1” bucket is associated with a different index (and corresponding directory **312**), the system **102** can uniquely identify them.

(124) Although illustrated as buckets “B1” and “B2,” it will be understood that the buckets (and/or corresponding sub-directories **314**) can be named in a variety of ways. In certain embodiments, the bucket (or sub-directory) names can include information about the bucket. For example, the bucket name can include the name of the index with which the bucket is associated, a time range of the

bucket, etc.

(125) As described herein, each bucket can have one or more files associated with it, including, but not limited to one or more raw machine data files, bucket summary files, filter files, inverted indexes (also referred to herein as high performance indexes or keyword indexes), permissions files, configuration files, etc. In the illustrated embodiment of FIG. 3B, the files associated with a particular bucket can be stored in the sub-directory corresponding to the particular bucket.

Accordingly, the files stored in the sub-directory **314A** can correspond to or be associated with bucket “B1,” of index “main,” the files stored in the sub-directory **314B** can correspond to or be associated with bucket “B2” of index “main,” and the files stored in the sub-directory **314C** can correspond to or be associated with bucket “B1” of index “test.”

(126) FIG. 3B further illustrates an expanded event data file **316C** showing an example of data that can be stored therein. In the illustrated embodiment, four events **320**, **322**, **324**, **326** of the machine data file **316C** are shown in four rows. Each event **320-326** includes machine data **330** and a timestamp **332**. The machine data **330** can correspond to the machine data received by the system **102**. For example, in the illustrated embodiment, the machine data **330** of events **320**, **322**, **324**, **326** corresponds to portions **302A**, **302B**, **302C**, **302D**, respectively, of the machine data **302** after it was processed by the indexing system **112**.

(127) Metadata **334-338** associated with the events **320-326** is also shown in the table **319**. In the illustrated embodiment, the metadata **334-338** includes information about a host **334**, source **336**, and sourcetype **338** associated with the events **320-326**. Any of the metadata can be extracted from the corresponding machine data, or supplied or defined by an entity, such as a user or computer system. The metadata fields **334-338** can become part of, stored with, or otherwise associated with the events **320-326**. In certain embodiments, the metadata **334-338** can be stored in a separate file of the sub-directory **314C** and associated with the machine data file **316C**. In some cases, while the timestamp **332** can be extracted from the raw data of each event, the values for the other metadata fields may be determined by the indexing system **112** based on information it receives pertaining to the host device **104** or data source of the data separate from the machine data.

(128) While certain default or user-defined metadata fields can be extracted from the machine data for indexing purposes, the machine data within an event can be maintained in its original condition. As such, in embodiments in which the portion of machine data included in an event is unprocessed or otherwise unaltered, it is referred to herein as a portion of raw machine data. For example, in the illustrated embodiment, the machine data of events **320-326** is identical to the portions of the machine data **302A-302D**, respectively, used to generate a particular event. Similarly, the entirety of the machine data **302** may be found across multiple events. As such, unless certain information needs to be removed for some reasons (e.g. extraneous information, confidential information), all the raw machine data contained in an event can be preserved and saved in its original form. Accordingly, the data store in which the event records are stored is sometimes referred to as a “raw record data store.” The raw record data store contains a record of the raw event data tagged with the various fields.

(129) In other embodiments, the portion of machine data in an event can be processed or otherwise altered relative to the machine data used to create the event. With reference to the machine data **304**, the machine data of a corresponding event (or events) may be modified such that only a portion of the machine data **304** is stored as one or more events. For example, in some cases, only machine data **304B** of the machine data **304** may be retained as one or more events or the machine data **304** may be altered to remove duplicate data, confidential information, etc.

(130) In FIG. 3B, the first three rows of the table **319** present events **320**, **322**, and **324** and are related to a server access log that records requests from multiple clients processed by a server, as indicated by entry of “access.log” in the source column **336**. In the example shown in FIG. 3B, each of the events **320-324** is associated with a discrete request made to the server by a client. The raw machine data generated by the server and extracted from a server access log can include the IP

address **1140** of the client, the user id **1141** of the person requesting the document, the time **1142** the server finished processing the request, the request line **1143** from the client, the status code **1144** returned by the server to the client, the size of the object **1145** returned to the client (in this case, the gif file requested by the client) and the time spent **1146** to serve the request in microseconds. In the illustrated embodiments of FIGS. **3A**, **3B**, all the raw machine data retrieved from the server access log is retained and stored as part of the corresponding events **320-324** in the file **316C**.

(131) Event **326** is associated with an entry in a server error log, as indicated by “error.log” in the source column **336** that records errors that the server encountered when processing a client request. Similar to the events related to the server access log, all the raw machine data in the error log file pertaining to event **326** can be preserved and stored as part of the event **326**.

(132) Saving minimally processed or unprocessed machine data in a data store associated with metadata fields in the manner similar to that shown in FIG. **3B** is advantageous because it allows search of all the machine data at search time instead of searching only previously specified and identified fields or field-value pairs. As mentioned above, because data structures used by various embodiments of the present disclosure maintain the underlying raw machine data and use a late-binding schema for searching the raw machines data, it enables a user to continue investigating and learn valuable insights about the raw data. In other words, the user is not compelled to know about all the fields of information that will be needed at data ingestion time. As a user learns more about the data in the events, the user can continue to refine the late-binding schema by defining new extraction rules, or modifying or deleting existing extraction rules used by the system.

(133) FIG. **3C** illustrates an embodiment of another file that can be included in one or more subdirectories **314** or buckets. Specifically, FIG. **3C** illustrates an exploded view of an embodiment of an inverted index **318B** in the sub-directory **314B**, associated with bucket “B2” of the index “main,” as well as an event reference array **340** associated with the inverted index **318B**.

(134) In some embodiments, the inverted indexes **318** can correspond to distinct time-series buckets. As such, each inverted index **318** can correspond to a particular range of time for an index. In the illustrated embodiment of FIG. **3C**, the inverted indexes **318A**, **318B** correspond to the buckets “B1” and “B2,” respectively, of the index “main,” and the inverted index **318C** corresponds to the bucket “B1” of the index “_test.” In some embodiments, an inverted index **318** can correspond to multiple time-series buckets (e.g., include information related to multiple buckets) or inverted indexes **318** can correspond to a single time-series bucket.

(135) Each inverted index **318** can include one or more entries, such as keyword (or token) entries **342** or field-value pair entries **344**. Furthermore, in certain embodiments, the inverted indexes **318** can include additional information, such as a time range **346** associated with the inverted index or an index identifier **348** identifying the index associated with the inverted index **318**. It will be understood that each inverted index **318** can include less or more information than depicted. For example, in some cases, the inverted indexes **318** may omit a time range **346** and/or index identifier **348**. In some such embodiments, the index associated with the inverted index **318** can be determined based on the location (e.g., directory **312**) of the inverted index **318** and/or the time range of the inverted index **318** can be determined based on the name of the sub-directory **314**.

(136) Token entries, such as token entries **342** illustrated in inverted index **318B**, can include a token **342A** (e.g., “error,” “itemID,” etc.) and event references **342B** indicative of events that include the token. For example, for the token “error,” the corresponding token entry includes the token “error” and an event reference, or unique identifier, for each event stored in the corresponding time-series bucket that includes the token “error.” In the illustrated embodiment of FIG. **3C**, the error token entry includes the identifiers 3, 5, 6, 8, 11, and 12 corresponding to events located in the bucket “B2” of the index “main.”

(137) In some cases, some token entries can be default entries, automatically determined entries, or user specified entries. In some embodiments, the indexing system **112** can identify each word or

string in an event as a distinct token and generate a token entry for the identified word or string. In some cases, the indexing system **112** can identify the beginning and ending of tokens based on punctuation, spaces, etc. In certain cases, the indexing system **112** can rely on user input or a configuration file to identify tokens for token entries **342**, etc. It will be understood that any combination of token entries can be included as a default, automatically determined, or included based on user-specified criteria.

(138) Similarly, field-value pair entries, such as field-value pair entries **344** shown in inverted index **318B**, can include a field-value pair **344A** and event references **344B** indicative of events that include a field value that corresponds to the field-value pair (or the field-value pair). For example, for a field-value pair sourcetype::sendmail, a field-value pair entry **344** can include the field-value pair “sourcetype::sendmail” and a unique identifier, or event reference, for each event stored in the corresponding time-series bucket that includes a sourcetype “sendmail.”

(139) In some cases, the field-value pair entries **344** can be default entries, automatically determined entries, or user specified entries. As a non-limiting example, the field-value pair entries for the fields “host,” “source,” and “sourcetype” can be included in the inverted indexes **318** as a default. As such, all of the inverted indexes **318** can include field-value pair entries for the fields “host,” “source,” and “sourcetype.” As yet another non-limiting example, the field-value pair entries for the field “IP_address” can be user specified and may only appear in the inverted index **318B** or the inverted indexes **318A**, **318B** of the index “_main” based on user-specified criteria. As another non-limiting example, as the indexing system **112** indexes the events, it can automatically identify field-value pairs and create field-value pair entries **344**. For example, based on the indexing system's **212** review of events, it can identify IP_address as a field in each event and add the IP address field-value pair entries to the inverted index **318B** (e.g., based on punctuation, like two keywords separated by an ‘=’ or ‘:’ etc.). It will be understood that any combination of field-value pair entries can be included as a default, automatically determined, or included based on user-specified criteria.

(140) With reference to the event reference array **340**, each unique identifier **350**, or event reference, can correspond to a unique event located in the time series bucket or machine data file **316B**. The same event reference can be located in multiple entries of an inverted index **318**. For example if an event has a sourcetype “splunkd,” host “www1” and token “warning,” then the unique identifier for the event can appear in the field-value pair entries **344** “sourcetype::splunkd” and “host::www1,” as well as the token entry “warning.” With reference to the illustrated embodiment of FIG. **3C** and the event that corresponds to the event reference **3**, the event reference **3** is found in the field-value pair entries **344** “host::hostA,” “source::sourceB,” “sourcetype::sourcetypeA,” and “IP_address::91.205.189.15” indicating that the event corresponding to the event references is from hostA, sourceB, of sourcetypeA, and includes “91.205.189.15” in the event data.

(141) For some fields, the unique identifier is located in only one field-value pair entry for a particular field. For example, the inverted index **318** may include four sourcetype field-value pair entries **344** corresponding to four different sourcetypes of the events stored in a bucket (e.g., sourcetypes: sendmail, splunkd, web_access, and web_service). Within those four sourcetype field-value pair entries, an identifier for a particular event may appear in only one of the field-value pair entries. With continued reference to the example illustrated embodiment of FIG. **3C**, since the event reference **7** appears in the field-value pair entry “sourcetype::sourcetypeA,” then it does not appear in the other field-value pair entries for the sourcetype field, including “sourcetype::sourcetypeB,” “sourcetype::sourcetypeC,” and “sourcetype::sourcetypeD.”

(142) The event references **350** can be used to locate the events in the corresponding bucket or machine data file **316**. For example, the inverted index **318B** can include, or be associated with, an event reference array **340**. The event reference array **340** can include an array entry **350** for each event reference in the inverted index **318B**. Each array entry **350** can include location information

352 of the event corresponding to the unique identifier (non-limiting example: seek address of the event, physical address, slice ID, etc.), a timestamp 354 associated with the event, or additional information regarding the event associated with the event reference, etc.

(143) For each token entry 342 or field-value pair entry 344, the event reference 342B, 344B, respectively, or unique identifiers can be listed in chronological order or the value of the event reference can be assigned based on chronological data, such as a timestamp associated with the event referenced by the event reference. For example, the event reference 1 in the illustrated embodiment of FIG. 3C can correspond to the first-in-time event for the bucket, and the event reference 12 can correspond to the last-in-time event for the bucket. However, the event references can be listed in any order, such as reverse chronological order, ascending order, descending order, or some other order (e.g., based on time received or added to the machine data file), etc. Further, the entries can be sorted. For example, the entries can be sorted alphabetically (collectively or within a particular group), by entry origin (e.g., default, automatically generated, user-specified, etc.), by entry type (e.g., field-value pair entry, token entry, etc.), or chronologically by when added to the inverted index, etc. In the illustrated embodiment of FIG. 3C, the entries are sorted first by entry type and then alphabetically.

(144) In some cases, inverted indexes 318 can decrease the search time of a query. For example, for a statistical query, by using the inverted index, the system 102 can avoid the computational overhead of parsing individual events in a machine data file 316. Instead, the system 102 can use the inverted index 318 separate from the raw record data store to generate responses to the received queries. Additional information is disclosed in U.S. Pat. No. 8,589,403, entitled “COMPRESSED JOURNALING IN EVENT TRACKING FILES FOR METADATA RECOVERY AND REPLICATION,” issued on 19 Nov. 2013; U.S. Pat. No. 8,412,696, entitled “REAL TIME SEARCHING AND REPORTING,” issued on 2 Apr. 2011; U.S. Pat. Nos. 8,589,375 and 8,589,432, both also entitled “REAL TIME SEARCHING AND REPORTING,” both issued on 19 Nov. 2013; U.S. Pat. No. 8,682,925, entitled “DISTRIBUTED HIGH PERFORMANCE ANALYTICS STORE,” issued on 25 Mar. 2014; U.S. Pat. No. 9,128,985, entitled “SUPPLEMENTING A HIGH PERFORMANCE ANALYTICS STORE WITH EVALUATION OF INDIVIDUAL EVENTS TO RESPOND TO AN EVENT QUERY,” issued on 8 Sep. 2015; and U.S. Pat. No. 9,990,386, entitled “GENERATING AND STORING SUMMARIZATION TABLES FOR SETS OF SEARCHABLE EVENTS,” issued on 5 Jun. 2018; each of which is hereby incorporated by reference in its entirety for all purposes, also include additional information regarding the creation and use of different embodiments of inverted indexes, including, but not limited to keyword indexes, high performance indexes and analytics stores, etc.

(145) 3.0. Query Processing and Execution

(146) FIG. 4A is a flow diagram illustrating an embodiment of a routine implemented by the query system 114 for executing a query. At block 402, the query system 114 receives a search query. As described herein, the query can be in the form of a pipelined (command) language or other) grouping of query language and includes filter criteria used to identify a set of data and processing criteria used to process the set of data.

(147) Examples of a pipelined (command) language may include Search Processing Language (SPL), any subsequent version of SPL (SPL2), or any pipelined languages of similar programming construct (e.g., generally referred to as “SPL”). As part of a pipelined search query, a sequence of operators is provided in which these operators may correspond to a command or a request to access and leverage the operability of a machine-learning (ML) component to perform a task (or tasks).

(148) Herein, the “ML component” may include, but is not limited or restricted to the following: a ML model (e.g., supervised ML algorithm trained with labeled data), a ML streaming algorithm (e.g., unsupervised ML algorithm trained with unlabeled data), a set of rules (e.g., one or more rules) that operate to perform the task or tasks, or a pipeline fragment (e.g., a published (recognized and operational) grouping of operators inclusive of the operator to access a ML model or a ML

streaming algorithm. Each of these operators, including an operator to access a ML component, is configured to process results provided by the prior operator in the series as input. Collectively, these operators may be referred to as an “operator sequence” or “SPL pipeline” while a subset of these operators inclusive of the operator associated with the ML component may be referred to as a “ML pipeline” or a “pipeline fragment.”

(149) At block **404**, the query system **114** processes the query. As part of processing the query, the query system **114** can determine whether the query was submitted by an authenticated user and/or review the query to determine that it is in a proper format for the data intake and query system **102**, has correct semantics and syntax, etc. In addition, the query system **114** can determine what, if any, configuration files or other configurations to use as part of the query.

(150) In addition as part of processing the query, the query system **114** can determine what portion(s) of the query to execute in a distributed manner (e.g., what to delegate to search nodes) and what portions of the query to execute in a non-distributed manner (e.g., what to execute on the search head). For the parts of the query that are to be executed in a distributed manner, the query system **114** can generate specific operators (e.g., commands, calls, etc.), for the components that are to execute the query. This may include generating subqueries, partial queries or different phases of the query for execution by different components of the query system **114**. In some cases, the query system **114** can use map-reduce techniques to determine how to map the data for the search and then reduce the data. Based on the map-reduce phases, the query system **114** can generate query operators for different components of the query system **114** or the SPL service component **506** of FIG. 5A as described below.

(151) As part of processing the query, the query system **114** can determine where to obtain the data. For example, in some cases, the data may reside on one or more indexing nodes or search nodes, as part of the storage system **116** or may reside in a shared storage system or a system external to the system **102**. In some cases, the query system **114** can determine what components to use to obtain and process the data. For example, the query system **114** can identify search nodes that are available for the query, etc.

(152) At block **406**, the query system **1206** distributes the determined portions or phases of the query to the appropriate components (e.g., search nodes). In some cases, the query system **1206** can use a catalog to determine which components to use to execute the query (e.g., which components include relevant data and/or are available, etc.).

(153) At block **408**, the components assigned to execute the query, execute the query. As mentioned, different components may execute different portions of the query. In some cases, multiple components (e.g., multiple search nodes) may execute respective portions of the query concurrently and communicate results of their portion of the query to another component (e.g., search head). As part of the identifying the set of data or applying the filter criteria, the components of the query system **114** can search for events that match the criteria specified in the query. These criteria can include matching keywords or specific values for certain fields. The searching operations at block **408** may use the late-binding schema to extract values for specified fields from events at the time the query is processed. In some embodiments, one or more rules for extracting field values may be specified as part of a sourcetype definition in a configuration file or in the query itself. In certain embodiments where search nodes are used to obtain the set of data, the search nodes can send the relevant events back to the search head, or use the events to determine a partial result, and send the partial result back to the search head.

(154) At block **410**, the query system **114** combines the partial results and/or events to produce a final result for the query. As mentioned, in some cases, combining the partial results and/or finalizing the results can include further processing the data according to the query. Such processing may entail joining different set of data, transforming the data, and/or performing one or more mathematical operations on the data, preparing the results for display, etc.

(155) In some examples, the results of the query are indicative of performance or security of the IT

environment and may help improve the performance of components in the IT environment. This final result may comprise different types of data depending on what the query requested. For example, the results can include a listing of matching events returned by the query, or some type of visualization of the data from the returned events. In another example, the final result can include one or more calculated values derived from the matching events.

(156) The results generated by the query system **114** can be returned to a client using different techniques. For example, one technique streams results or relevant events back to a client in real-time as they are identified. Another technique waits to report the results to the client until a complete set of results (which may include a set of relevant events or a result based on relevant events) is ready to return to the client. Yet another technique streams interim results or relevant events back to the client in real-time until a complete set of results is ready, and then returns the complete set of results to the client. In another technique, certain results are stored as “search jobs” and the client may retrieve the results by referring to the search jobs.

(157) The query system **114** can also perform various operations to make the search more efficient. For example, before the query system **114** begins execution of a query, it can determine a time range for the query and a set of common keywords that all matching events include. The query system **114** may then use these parameters to obtain a superset of the eventual results. Then, during a filtering stage, the query system **114** can perform field-extraction operations on the superset to produce a reduced set of search results. This speeds up queries, which may be particularly helpful for queries that are performed on a periodic basis. In some cases, to make the search more efficient, the query system **114** can use information known about certain data sets that are part of the query to filter other data sets. For example, if an early part of the query includes instructions to obtain data with a particular field, but later commands of the query do not rely on the data with that particular field, the query system **114** can omit the superfluous part of the query from execution.

(158) 3.1. Pipelined Search Language

(159) Various embodiments of the present disclosure can be implemented using, or in conjunction with, a pipelined language. A pipelined language is a language in which a set of inputs or data is operated on by a first operator (e.g., command, function or system call directed to a ML component, etc.) in a sequence of commands, and then subsequent operators in the order they are arranged in the sequence. Such operators can include any type of functionality for operating on data, such as retrieving, searching, filtering, aggregating, processing, transmitting, and the like. As described herein, a query can thus be formulated in a pipelined (command) language and include any number of ordered or unordered operators for operating on data.

(160) Search Processing Language (SPL) is an example of a pipelined (command) language in which a set of inputs or data is operated on by any number of operators in a particular sequence. A sequence of operators, referred to as an or “operator sequence,” can be formulated such that the order in which the operators are arranged defines the order in which the operators are applied to a set of data or the results of an earlier executed operator. For example, a first operator in an operator sequence normally corresponds to a command that can include filter criteria used to search or filter for specific data. The results of the first operator can then be passed to another operator listed later in the operator sequence for further processing.

(161) In various embodiments, a query can be formulated as an operator sequence defined in a command line of a search UI. In some embodiments, a query can be formulated as a sequence of SPL operators. Some or all of the SPL operators in the sequence of SPL operators can be separated from one another by a pipe symbol “|.” In such embodiments, a set of data, such as a set of events, can be operated on by a first SPL operator in the sequence, and then a subsequent SPL operator following a pipe symbol “|” after the first SPL operator operates on the results produced by the first SPL operator or other set of data, and so on for any additional SPL operators in the sequence. As such, a query formulated using SPL comprises a series of consecutive operators that are delimited by pipe “|” characters. The pipe character indicates to the system that the output or result of one

operator (to the left of the pipe) should be used as the input for one of the subsequent operators (to the right of the pipe). This enables formulation of queries defined by a pipeline of sequenced commands that refines or enhances the data at each step along the pipeline until the desired results are attained. Accordingly, various embodiments described herein can be implemented with Search Processing Language (SPL) used in conjunction with the SPLUNK® ENTERPRISE system.

(162) While a query can be formulated in many ways, a query can start with a search command and one or more corresponding search terms or filter criteria at the beginning of the pipeline. Such search terms or filter criteria can include any combination of keywords, phrases, times, dates, Boolean expressions, fieldname-field value pairs, etc. that specify which results should be obtained from different locations. The results can then be passed as inputs into subsequent operators in a sequence of operators by using, for example, a pipe character. The subsequent operators in a sequence can include directives for additional processing of the results once it has been obtained from one or more indexes. For example, operators may be used to filter unwanted information out of the results, extract more information, evaluate field values, calculate statistics, reorder the results, create an alert, create summary of the results, or perform some type of aggregation function. In some embodiments, the summary can include a graph, chart, metric, or other visualization of the data. An aggregation function can include analysis or calculations to return an aggregate value, such as an average value, a sum, a maximum value, a root mean square, statistical values, and the like.

(163) Due to its flexible nature, use of a pipelined (command) language in various embodiments is advantageous because it can perform “filtering” as well as “processing” functions. In other words, a single query can include a search command and search term expressions, as well as data-analysis expressions. For example, a command at the beginning of a query can perform a “filtering” step by retrieving a set of data based on a condition (e.g., records associated with server response times of less than 1 microsecond). The results of the filtering step can then be passed to a subsequent operator in the pipeline that performs a “processing” step (e.g. calculation of an aggregate value related to the filtered events such as the average response time of servers with response times of less than 1 microsecond). Furthermore, the search command can allow events to be filtered by keyword as well as field criteria. For example, a search command can filter events based on the word “warning” or filter events based on a field value “10.0.1.2” associated with a field “clientip.”

(164) The results obtained or generated in response to an operator in a query can be considered a set of results data. The set of results data can be passed from one operator to another in any data format. In one embodiment, the set of result data can be in the form of a dynamically created table. Each operator in a particular query can redefine the shape of the table. In some implementations, an event retrieved from an index in response to a query can be considered a row with a column for each field value. Columns can contain basic information about the data and/or data that has been dynamically extracted at search time.

(165) FIG. 4B provides a visual representation of the manner in which a pipelined language or query can operate in accordance with the disclosed embodiments. The query **430** can be input by the user and submitted to the query system **114**. In the illustrated embodiment, the query **430** comprises filter criteria **430A**, followed by two operators represented as commands **430B**, **430C** (namely, Command1 and Command2). Disk **422** represents data as it is stored in a data store to be searched. For example, disk **422** can represent a portion of the storage system **116** or some other data store that can be searched by the query system **114**. Individual rows of can represent different events and columns can represent different fields for the different events. In some cases, these fields can include raw machine data, host, source, and sourcetype.

(166) At block **440**, the query system **114** uses the filter criteria **430A** (e.g., “sourcetype=syslog ERROR”) to filter events stored on the disk **422** to generate an intermediate results table **424**. Given the semantics of the query **430** and order of the commands, the query system **114** can execute the filter criteria **430A** portion of the query **430** before executing Command1 or

Command2.

(167) Rows in the table **424** may represent individual records, where each record corresponds to an event in the disk **422** that satisfied the filter criteria. Columns in the table **424** may correspond to different fields of an event or record, such as “user,” “count,” “percentage,” “timestamp,” or the raw machine data of an event, etc. Notably, the fields in the intermediate results table **424** may differ from the fields of the events on the disk **422**. In some cases, this may be due to the late binding schema described herein that can be used to extract field values at search time. Thus, some of the fields in table **424** may not have existed in the events on disk **422**.

(168) Illustratively, the intermediate results table **424** has fewer rows than what is shown in the disk **422** because only a subset of events retrieved from the disk **422** matched the filter criteria **430A** “sourcetype=syslog ERROR.” In some embodiments, instead of searching individual events or raw machine data, the set of events in the intermediate results table **424** may be generated by a call to a pre-existing inverted index.

(169) At block **442**, the query system **114** processes the events of the first intermediate results table **424** to generate the second intermediate results table **426**. With reference to the query **430**, the query system **114** processes the events of the first intermediate results table **424** to identify the top users according to Command1. This processing may include determining a field value for the field “user” for each record in the intermediate results table **424**, counting the number of unique instances of each “user” field value (e.g., number of users with the name David, John, Julie, etc.) within the intermediate results table **424**, ordering the results from largest to smallest based on the count, and then keeping only the top 10 results (e.g., keep an identification of the top 10 most common users). Accordingly, each row of table **426** can represent a record that includes a unique field value for the field “user,” and each column can represent a field for that record, such as fields “user,” “count,” and “percentage.”

(170) At block **444**, the query system **114** processes the second intermediate results table **426** to generate the final results table **428**. With reference to query **430**, the query system **114** applies the command “fields-present” to the second intermediate results table **426** to generate the final results table **428**. As shown, the command “fields-present” of the query **430** results in one less column, which may represent that a field was removed during processing. For example, the query system **114** may have determined that the field “percentage” was unnecessary for displaying the results based on the Command2. In such a scenario, each record of the final results table **428** would include a field “user,” and “count.” Further, the records in the table **428** would be ordered from largest count to smallest count based on the query commands.

(171) It will be understood that the final results table **428** can be a third intermediate results table, which can be pipelined to another stage where further filtering or processing of the data can be performed, e.g., preparing the data for display purposes, filtering the data based on a condition, performing a mathematical calculation with the data, etc. In different embodiments, other query languages, such as the Structured Query Language (“SQL”), can be used to create a query.

(172) 3.2. Field Extraction

(173) As described herein, extraction rules can be used to extract field-value pairs or field values from data. An extraction rule can comprise one or more regex rules that specify how to extract values for the field corresponding to the extraction rule. In addition to specifying how to extract field values, the extraction rules may also include instructions for deriving a field value by performing a function on a character string or value retrieved by the extraction rule. For example, an extraction rule may truncate a character string or convert the character string into a different data format. Extraction rules can be used to extract one or more values for a field from events by parsing the portions of machine data in the events and examining the data for one or more patterns of characters, numbers, delimiters, etc., that indicate where the field begins and, optionally, ends. In certain embodiments, extraction rules can be stored in one or more configuration files. In some cases, a query itself can specify one or more extraction rules.

(174) In some cases, extraction rules can be applied at data ingest by the intake system **110** and/or indexing system **112**. For example, the intake system **110** and indexing system **112** can apply extraction rules to ingested data and/or events generated from the ingested data and store results in an inverted index.

(175) The system **102** advantageously allows for search time field extraction. In other words, fields can be extracted from the event data at search time using late-binding schema as opposed to at data ingestion time, which was a major limitation of the prior art systems. Accordingly, extraction rules can be applied at search time by the query system **114**. The query system can apply extraction rules to events retrieved from the storage system **116** or data received from sources external to the system **102**. Extraction rules can be applied to all the events in the storage system **116** or to a subset of the events that have been filtered based on some filter criteria (e.g., event timestamp values, etc.).

(176) FIG. 4C is a block diagram illustrating an embodiment of the table **319** showing events **320-326**, described previously with reference to FIG. 3B. As described herein, the table **319** is for illustrative purposes, and the events **320-326** may be stored in a variety of formats in an event data file **316** or raw record data store. Further, it will be understood that the event data file **316** or raw record data store can store millions of events. FIG. 4C also illustrates an embodiment of a search bar **450** for entering a query and a configuration file **452** that includes various extraction rules that can be applied to the events **320-326**.

(177) As a non-limiting example, if a user inputs a query into search bar **450** that includes only keywords (also known as “tokens”), e.g., the keyword “error” or “warning,” the query system **114** can search for those keywords directly in the events **320-326** stored in the raw record data store.

(178) As described herein, the indexing system **112** can optionally generate and use an inverted index with keyword entries to facilitate fast keyword searching for event data. If a user searches for a keyword that is not included in the inverted index, the query system **114** may nevertheless be able to retrieve the events by searching the event data for the keyword in the event data file **316** or raw record data store directly. For example, if a user searches for the keyword “eva,” and the name “eva” has not been indexed at search time, the query system **114** can search the events **320-326** directly and return the first event **320**. In the case where the keyword has been indexed, the inverted index can include a reference pointer that will allow for a more efficient retrieval of the event data from the data store. If the keyword has not been indexed, the query system **114** can search through the events in the event data file to service the search.

(179) In many cases, a query include fields. The term “field” refers to a location in the event data containing one or more values for a specific data item. Often, a field is a value with a fixed, delimited position on a line, or a name and value pair, where there is a single value to each field name. A field can also be multivalued, that is, it can appear more than once in an event and have a different value for each appearance, e.g., email address fields. Fields are searchable by the field name or field name-value pairs. Some examples of fields are “clientip” for IP addresses accessing a web server, or the “From” and “To” fields in email addresses.

(180) By way of further example, consider the query, “status=404.” This search query finds events with “status” fields that have a value of “404.” When the search is run, the query system **114** does not look for events with any other “status” value. It also does not look for events containing other fields that share “404” as a value. As a result, the search returns a set of results that are more focused than if “404” had been used in the search string as part of a keyword search. Note also that fields can appear in events as “key=value” pairs such as “user_name=Bob.” But in most cases, field values appear in fixed, delimited positions without identifying keys. For example, the data store may contain events where the “user_name” value always appears by itself after the timestamp as illustrated by the following string: “Nov. 15 09:33:22 evaemerson.”

(181) FIG. 4C illustrates the manner in which configuration files may be used to configure custom fields at search time in accordance with the disclosed embodiments. In response to receiving a

query, the query system **114** determines if the query references a “field.” For example, a query may request a list of events where the “clientip” field equals “127.0.0.1.” If the query itself does not specify an extraction rule and if the field is not an indexed metadata field, e.g., time, host, source, sourcetype, etc., then in order to determine an extraction rule, the query system **114** may, in one or more embodiments, locate configuration file **452** during the execution of the query.

(182) Configuration file **452** may contain extraction rules for various fields, e.g., the “clientip” field. The extraction rules may be inserted into the configuration file **452** in a variety of ways. In some embodiments, the extraction rules can comprise regular expression rules that are manually entered in by the user.

(183) In one or more embodiments, as noted above, a field extractor may be configured to automatically generate extraction rules for certain field values in the events when the events are being created, indexed, or stored, or possibly at a later time. In one embodiment, a user may be able to dynamically create custom fields by highlighting portions of a sample event that should be extracted as fields using a graphical user interface. The system can then generate a regular expression that extracts those fields from similar events and store the regular expression as an extraction rule for the associated field in the configuration file **452**.

(184) In some embodiments, the indexing system **112** can automatically discover certain custom fields at index time and the regular expressions for those fields will be automatically generated at index time and stored as part of extraction rules in configuration file **452**. For example, fields that appear in the event data as “key=value” pairs may be automatically extracted as part of an automatic field discovery process. Note that there may be several other ways of adding field definitions to configuration files in addition to the methods discussed herein.

(185) Events from heterogeneous sources that are stored in the storage system **116** may contain the same fields in different locations due to discrepancies in the format of the data generated by the various sources. For example, event **326** also contains a “clientip” field, however, the “clientip” field is in a different format from events **320**, **322**, and **324**. Furthermore, certain events may not contain a particular field at all. To address the discrepancies in the format and content of the different types of events, the configuration file **452** can specify the set of events to which an extraction rule applies. For example, extraction rule **454** specifies that it is to be used with events having a sourcetype “access_combined,” and extraction rule **456** specifies that it is to be used with events having a sourcetype “apache_error.” Other extraction rules shown in configuration file **452** specify a set or type of events to which they apply. In addition, the extraction rules shown in configuration file **452** include a regular expression for parsing the identified set of events to determine the corresponding field value. Accordingly, each extraction rule may pertain to only a particular type of event. Accordingly, if a particular field, e.g., “clientip” occurs in multiple types of events, each of those types of events can have its own corresponding extraction rule in the configuration file **452** and each of the extraction rules would comprise a different regular expression to parse out the associated field value. In some cases, the sets of events are grouped by sourcetype because events generated by a particular source can have the same format.

(186) The field extraction rules stored in configuration file **452** can be used to perform search-time field extractions. For example, for a query that requests a list of events with sourcetype “access_combined” where the “clientip” field equals “127.0.0.1,” the query system **114** can locate the configuration file **452** to retrieve extraction rule **454** that allows it to extract values associated with the “clientip” field from the events where the sourcetype is “access_combined” (e.g., events **320-324**). After the “clientip” field has been extracted from the events **320**, **322**, **324**, the query system **114** can then apply the field criteria by performing a compare operation to filter out events where the “clientip” field does not equal “127.0.0.1.” In the example shown in FIG. **4C**, the events **320** and **322** would be returned in response to the user query. In this manner, the query system **114** can service queries with filter criteria containing field criteria and/or keyword criteria.

(187) It should also be noted that any events filtered by performing a search-time field extraction

using a configuration file **452** can be further processed by directing the results of the filtering step to a processing step using a pipelined search language. Using the prior example, a user can pipeline the results of the compare step to an aggregate function by asking the query system **114** to count the number of events where the “clientip” field equals “127.0.0.1.”

(188) By providing the field definitions for the queried fields at search time, the configuration file **452** allows the event data file or raw record data store to be field searchable. In other words, the raw record data store can be searched using keywords as well as fields, wherein the fields are searchable name/value pairings that can distinguish one event from another event and can be defined in configuration file **452** using extraction rules. In comparison to a search containing field names, a keyword search may result in a search of the event data directly without the use of a configuration file.

(189) Further, the ability to add schema to the configuration file **452** at search time results in increased efficiency and flexibility. A user can create new fields at search time and simply add field definitions to the configuration file **452**. As a user learns more about the data in the events, the user can continue to refine the late-binding schema by adding new fields, deleting fields, or modifying the field extraction rules in the configuration file for use the next time the schema is used by the system **102**. Because the system **102** maintains the underlying raw data and uses late-binding schema for searching the raw data, it enables a user to continue investigating and learn valuable insights about the raw data long after data ingestion time. Similarly, multiple field definitions can be added to the configuration file to capture the same field across events generated by different sources or sourcetypes. This allows the system **102** to search and correlate data across heterogeneous sources flexibly and efficiently.

(190) 3.3. Data Models

(191) The system **102** can use one or more data models to search and/or better understand data. A data model is a hierarchically structured search-time mapping of semantic knowledge about one or more datasets. It encodes the domain knowledge used to build a variety of specialized searches of those datasets. Those searches, in turn, can be used to generate reports. Additional information regarding data models, their creation and their use is described in U.S. Pat. Nos. 8,788,525 and 8,788,526, both entitled “DATA MODEL FOR MACHINE DATA FOR SEMANTIC SEARCH,” both issued on 22 Jul. 2014; U.S. Pat. No. 8,983,994, entitled “GENERATION OF A DATA MODEL FOR SEARCHING MACHINE DATA,” issued on 17 Mar. 2015; U.S. Pat. No. 9,128,980, entitled “GENERATION OF A DATA MODEL APPLIED TO QUERIES,” issued on 8 Sep. 2015; U.S. Pat. No. 9,589,012, entitled “GENERATION OF A DATA MODEL APPLIED TO OBJECT QUERIES,” issued on 7 Mar. 2017; each of which is hereby incorporated by reference in its entirety for all purposes. Building reports using a report generation interface is further explained in U.S. patent application Ser. No. 14/503,335, entitled “GENERATING REPORTS FROM UNSTRUCTURED DATA,” filed on 30 Sep. 2014, and which is hereby incorporated by reference in its entirety for all purposes.

(192) 3.4. Acceleration Techniques

(193) The above-described system provides significant flexibility by enabling a user to analyze massive quantities of minimally-processed data “on the fly” at search time using a late-binding schema, instead of storing pre-specified portions of the data in a database at ingestion time. This flexibility enables a user to see valuable insights, correlate data, and perform subsequent queries to examine interesting aspects of the data that may not have been apparent at ingestion time.

(194) Performing extraction and analysis operations at search time can involve a large amount of data and require a large number of computational operations, which can cause delays in processing the queries. In some embodiments, the system **102** can employ a number of unique acceleration techniques to speed up analysis operations performed at search time. These techniques include: performing search operations in parallel using multiple components of the query system **114**, using an inverted index **118**, and accelerating the process of generating reports.

(195) To facilitate faster query processing, a query can be structured such that multiple components of the query system **114** (e.g., search nodes) perform the query in parallel, while aggregation of search results from the multiple components is performed at a particular component (e.g., search head). For example, consider a scenario in which a user enters the query “Search “error” | stats count BY host.” The query system **114** can identify two phases for the query, including: (1) subtasks (e.g., data retrieval or simple filtering) that may be performed in parallel by multiple components, such as search nodes, and (2) a search results aggregation operation to be executed by one component, such as the search head, when the results are ultimately collected from the search nodes.

(196) Based on this determination, the query system **114** can generate commands to be executed in parallel by the search nodes, with each search node applying the generated commands to a subset of the data to be searched. In this example, the query system **114** generates and then distributes the following commands to the individual search nodes: “Search “error” | prestats count BY host.” In this example, the “prestats” command can indicate that individual search nodes are processing a subset of the data and are responsible for producing partial results and sending them to the search head. After the search nodes return the results to the search head, the search head aggregates the received results to form a single search result set. By executing the query in this manner, the system effectively distributes the computational operations across the search nodes while reducing data transfers. It will be understood that the query system **114** can employ a variety of techniques to use distributed components to execute a query. In some embodiments, the query system **114** can use distributed components for only mapping functions of a query (e.g., gather data, applying filter criteria, etc.). In certain embodiments, the query system **114** can use distributed components for mapping and reducing functions (e.g., joining data, combining data, reducing data, etc.) of a query.

(197) 4.0. Example Use Cases

(198) The system **102** provides various schemas, dashboards, and visualizations that simplify developers' tasks to create applications with additional capabilities, including but not limited to security, data center monitoring, IT service monitoring, and client/customer insights.

(199) 4.1. Security Features

(200) An embodiment of an enterprise security application is as SPLUNK® ENTERPRISE SECURITY, which performs monitoring and alerting operations and includes analytics to facilitate identifying both known and unknown security threats based on large volumes of data stored by the system **102**. The enterprise security application provides the security practitioner with visibility into security-relevant threats found in the enterprise infrastructure by capturing, monitoring, and reporting on data from enterprise security devices, systems, and applications. Through the use of the system **102** searching and reporting capabilities, the enterprise security application provides a top-down and bottom-up view of an organization's security posture.

(201) Additional information related to various security applications is described in U.S. application Ser. No. 16/512,899; U.S. Pat. No. 8,826,434, entitled “SECURITY THREAT DETECTION BASED ON INDICATIONS IN BIG DATA OF ACCESS TO NEWLY REGISTERED DOMAINS,” issued on 2 Sep. 2014; U.S. Pat. No. 9,215,240, entitled “INVESTIGATIVE AND DYNAMIC DETECTION OF POTENTIAL SECURITY-THREAT INDICATORS FROM EVENTS IN BIG DATA,” issued on 15 Dec. 2015; U.S. Pat. No. 9,173,801, entitled “GRAPHIC DISPLAY OF SECURITY THREATS BASED ON INDICATIONS OF ACCESS TO NEWLY REGISTERED DOMAINS,” issued on 3 Nov. 2015; U.S. Pat. No. 9,248,068, entitled “SECURITY THREAT DETECTION OF NEWLY REGISTERED DOMAINS,” issued on 2 Feb. 2016; U.S. Pat. No. 9,426,172, entitled “SECURITY THREAT DETECTION USING DOMAIN NAME ACCESSES,” issued on 23 Aug. 2016; U.S. Pat. No. 9,432,396, entitled “SECURITY THREAT DETECTION USING DOMAIN NAME REGISTRATIONS,” issued on 30 Aug. 2016; and U.S. App. Pub. No. 2013/0318236, entitled “KEY INDICATORS VIEW,” filed on 31 Jul. 2013; each of which is hereby incorporated by

reference in its entirety for all purposes.

(202) 4.3. IT Service Monitoring

(203) An embodiment of an IT monitoring application is SPLUNK® IT SERVICE INTELLIGENCE™, which performs monitoring and alerting operations. The IT monitoring application also includes analytics to help an analyst diagnose the root cause of performance problems based on large volumes of data stored by the system **102** as correlated to the various services an IT organization provides (a service-centric view). This differs significantly from conventional IT monitoring systems that lack the infrastructure to effectively store and analyze large volumes of service-related events. Traditional service monitoring systems typically use fixed schemas to extract data from pre-defined fields at data ingestion time, wherein the extracted data is typically stored in a relational database. This data extraction process and associated reduction in data content that occurs at data ingestion time inevitably hampers future investigations, when all of the original data may be needed to determine the root cause of or contributing factors to a service issue.

(204) In contrast, an IT monitoring application system stores large volumes of minimally-processed service-related data at ingestion time for later retrieval and analysis at search time, to perform regular monitoring, or to investigate a service issue. To facilitate this data retrieval process, the IT monitoring application enables a user to define an IT operations infrastructure from the perspective of the services it provides. In this service-centric approach, a service such as corporate e-mail may be defined in terms of the entities employed to provide the service, such as host machines and network devices. Each entity is defined to include information for identifying all of the events that pertains to the entity, whether produced by the entity itself or by another machine, and considering the many various ways the entity may be identified in machine data (such as by a URL, an IP address, or machine name). The service and entity definitions can organize events around a service so that all of the events pertaining to that service can be easily identified. This capability provides a foundation for the implementation of Key Performance Indicators.

(205) Additional disclosure regarding IT Service Monitoring is described in U.S. application Ser. No. 16/512,899, incorporated by reference herein in its entirety.

(206) 4.4. Client or Customer Insights

(207) As described herein, the system **102** can receive heterogeneous data from disparate systems. In some cases, the data from the disparate systems may be related and correlating the data can result in insights into client or customer interactions with various systems of a vendor. To aid in the correlation of data across different systems, multiple field definitions can be added to one or more configuration files to capture the same field or data across events generated by different sources or sourcetypes. This can enable the system **102** to search and correlate data across heterogeneous sources flexibly and efficiently.

(208) As a non-limiting example and with reference to FIG. 4D, consider a scenario in which a common customer identifier is found among log data received from three disparate data sources. In this example, a user submits an order for merchandise using a vendor's shopping application program **460** running on the user's system. In this example, the order was not delivered to the vendor's server due to a resource exception at the destination server that is detected by the middleware code **462**. The user then sends a message to the customer support server **464** to complain about the order failing to complete. The three systems **460**, **462**, **464** are disparate systems that do not have a common logging format. The shopping application program **460** sends log data **466** to the system **102** in one format, the middleware code **462** sends error log data **468** in a second format, and the support server **464** sends log data **470** in a third format.

(209) Using the log data received at the system **102** from the three systems **460**, **462**, **464**, the vendor can uniquely obtain an insight into user activity, user experience, and system behavior. The system **102** allows the vendor's administrator to search the log data from the three systems **460**, **462**, **464**, thereby obtaining correlated information, such as the order number and corresponding

customer ID number of the person placing the order. The system **102** also allows the administrator to see a visualization of related events via a user interface. The administrator can query the system **102** for customer ID field value matches across the log data from the three systems **460, 462, 464** that are stored in the storage system **116**. While the customer ID field value exists in the data gathered from the three systems **460, 462, 464**, it may be located in different areas of the data given differences in the architecture of the systems. The query system **114** obtains events from the storage system **116** related to the three systems **460, 462, 464**. The query system **114** then applies extraction rules to the events in order to extract field values for the field “customer ID” that it can correlate. As described herein, the query system **114** may apply a different extraction rule to each set of events from each system when the event format differs among systems. In this example, a user interface can display to the administrator the events corresponding to the common customer ID field values **472, 474, and 476**, thereby providing the administrator with insight into a customer's experience. The system **102** can provide additional user interfaces and reports to aid a user in analyzing the data associated with the customer.

(210) 5.0. System Architecture Specific Description

(211) Modern data scientists have access to several tools or platforms for developing and executing software. Generally, data scientists seek to develop software using a platform that provides support for multiple programming languages, efficiently handles reading of structured and unstructured data, places no restrictions on the use or loading of external libraries, allows for collaboration among multiple data scientists and utilizes scalable computing resources configured to execute aspects of the developed software. However, currently there are no such available platforms for use by data scientists providing each of these desired advantages.

(212) As a result, aspects of the disclosure describe a data science platform (or a “unified analytics platform”) that provides such advantages and additionally provides native support for various query languages including SPL, as discussed above. In particular, embodiments disclosed herein include a data science platform that provides seamless integration into applications that receive programming code, which may include SPL queries. The platform may receive these SPL queries when they are to be executed. In some embodiments, the platform parses a received query to determine a data source, a data sink and operators comprising the query, optionally alters the ordering of and/or consolidates operators in order to improve the execution efficiency of the query, converts the query (or altered query) to a directed acyclic graph (DAG), and provides the DAG to a distributed processing engine for execution.

(213) In some embodiments, logic of the data science platform, e.g., a kernel logic **505** of the data science platform **504** of FIGS. 5A-5B, may directly interface with a client-based interactive web application such as a JUPYTER® notebook, where the kernel logic may be configured to provide both syntax highlighting, an autocomplete functionality (e.g., to provide programming syntax suggestions), and perform conversions between SPL, Python, R, Scala, etc.

(214) Additionally, some embodiments of the platform described in further detail below include receipt of third-party code to be stored by the platform. Third-party code (or “untrusted code” or “unsecured code”) may refer to code written, or otherwise provided, by a data scientist that is to be utilized by the platform in performing operations of parsing, altering and/or converting to a DAG. In some embodiments, the third-party code is a machine learning model or representation thereof. Embodiments of the platform include specific components (or “logic”) configured to analyze the third-party code thereby ensuring the code is secure, meaning the code may be executed without anomalous, malicious or undesired results such as attempts to access certain data, alter permissions of various applications, cause an unexpected termination of the execution process, etc. The components of the platform specifically configured to secure the third-party code may verify the schema of the third-party code and further validate the code. In instances when the third-party code is a machine learning model, the input schema of the machine learning model is verified as being consistent in accordance with the schema deduced from the training of the model and validate the

model through deserialization of the machine learning to a computational graph representing the model. For example, validation may include deserialization of the machine learning model into one or more computational graphs.

(215) The embodiments of the data science platform described herein include a platform comprising several components and one or more non-transitory, computer-readable medium storage components that interact with a front-end application (e.g., a software programming application) and also interacts with a distributed processing engine. In particular, the platform interacts with the front-end application for two purposes: (1) to receive queries for processing; and (2) to receive third-party code for validation and storage by the platform. Additionally, the platform interacts with the distributed processing engine to provide the engine with a DAG representing a query. As a result, the platform enables a data scientist to develop software code, including a query, and have that query executed by a distributed processing engine, which is auto-scaling and separate from the storage of data on which the query is executed.

(216) Thus, the data science platform described herein provides a distinct advantage over current query execution platforms by decoupling the execution of the query and the storage of the data on which the query is executed. For instance, current query executions platforms include a compute node that executes a query by retrieving all data indicated within the query and then executes each operator within the query (i.e., in a pipelined manner as discussed above). Such coupling of the compute (execution) and storage is often inefficient, especially when the amount of data retrieved as an initial step in execution is large (e.g., a table containing thousands of columns and thousands of rows). Therefore, decoupling the compute from the storage provides significant computational benefits (e.g., resource saving) by decreasing the amount of time resources are utilized while executing a single query. Additionally, as will be discussed below, execution of the query by a distributed processing engine allows for the ability to parallelize data reads, which obviates traditional bottlenecks during query execution.

(217) Additionally, the data science platform disclosed provides technical advantages over the prior art while providing improvements to the technical field of query execution by efficiently retrieving and parsing unstructured data. As is typical in the art, data that is to be retrieved for the purpose of utilization during execution of a search query is structured at the front-end, e.g., at the time of storage. However, there are numerous inefficiencies with this including, but not limited or restricted to, spending time predefining a structure and utilization of computing resources to structure the data upon receipt prior to storage. In contrast, embodiments of the disclosed data science platform include the ability for a data scientist to utilize a technology add-on (“TA”) directly in a query and/or the functionality to determine or infer a structure of unstructured data. A TA may be a logic module or component that includes instructions detailing how data from a particular source is to be retrieved. The information provided by a TA may be referred to as a search-time knowledge map that provides instructions for normalizing the data. In some embodiments, the data science platform may perform an initial “preliminary query” that samples unstructured data to determine or infer an organization of the data in order to efficiently extract this data in accordance with a query.

(218) Further, certain embodiments of the data science platform described herein include components specifically configured to alter a received query to improve the efficiency of the execution of the query. For instance, the components may reorder and/or consolidate operators of the received query thereby generating an altered query. For example, certain operators may be consolidated such that a filtering of data is performed prior to or concurrently with (at least partially overlapping in time) the retrieval of the data indicated by the query. Thus, the amount of data initially retrieved may be greatly reduced, which provides further resource saving benefits. As one example, which will be discussed in further detail below, altering of the query may include an analysis of the operators comprising the query to determine alterations that may be made to reduce the amount of retrieved data by applying a filter and/or a projection on an operator that retrieves

data (“data retrieval operator,” e.g., read_json (“ . . . ”)). Specifically, operators subsequent to the data retrieval operator may be analyzed to determine which rows and/or columns of a data set are required by the query and a filter and/or a projection may be applied to the data retrieval operator such that only the data required by the operators comprising the query are retrieved.

(219) 5.0.1. Architecture Over View

(220) Referring now to FIG. 5A, a block diagram of an embodiment of a data processing environment **500** including a data science platform **504** is shown in accordance with some embodiments. The data processing environment **500** includes a virtual environment **502** and the data science platform **504** that includes a kernel logic **505** (operating within the virtual environment **502**), a SPL service component **506**, a model service component **508** and an operator storage **510**, an auto-scaling, clustered, distributed processing engine **512**, a plurality of data intake and query systems **108.sub.1-108.sub.i** (where $i \geq 1$), a plurality of data sources **514.sub.1-514.sub.j** (where $j \geq 1$), and a plurality of data lakes **516.sub.1-516.sub.k** (where $k \geq 1$) (or other data reservoirs).

(221) The virtual environment **502** may include one or more instances of an application configured to receive software code (“a programming application”), where the one or more instances are communicatively coupled to the data science platform **504** such that the SPL service component **506** and the model service component **508** may obtain software code and/or queries from the one or more instances. The software code may include a machine learning model or a representation thereof. The queries may be a software code formatted in a specific syntax using a known query language. In some embodiments, a query may comprise a pipelined search query, and in some specific embodiments, the query may be a SPL query (discussed above and any version thereof). In some embodiments, the kernel logic **505** may be the interface between the instances of the application (e.g., a JUPYTER® notebook instance) and the data science platform **504**. In particular, the kernel logic **505** may be configured to provide both syntax highlighting, an autocomplete functionality (e.g., to provide programming syntax suggestions), and perform conversions between SPL, Python, R, Scala, etc. For example, when a programming language other than SPL (e.g., Python, R, Scala, etc.) is utilized within a SPL query, the kernel logic **505** may convert the syntax of that programming language to SPL prior to providing the SPL query to the data science platform **504**. Additionally, in some embodiments, the kernel logic **505** may be configured to and responsible for converting a machine learning model to an Open Neural Network Exchange (ONNX) model, as discussed below. Further, the kernel logic **505** may be configured to receive a handle providing access to results of a query and/or a status of a query that is displayed within the JUPTYER® notebook interface. Although the kernel logic **505** is illustrated as a single instance in FIGS. 5A-5B, in some embodiments, an instance may be provided for and operating within each web application instance (e.g., JUPYTER® notebook instance) or each sandboxed environment **518.sub.1-518.sub.L** (as seen in FIG. 5B).

(222) In yet additional embodiments, the kernel logic **505** is configured to dispatch queries to the data science platform **504** and collect results of the execution of the queries into data structures that are configured to be provided to commands or operators in downstream manipulations (e.g., written in a programming language other than SPL such as R, Python, Scala, etc.). In yet other embodiments, the kernel logic **505** may be configured to parameterize SPL queries and cache results of their execution.

(223) The components of the data science platform **504** may be configured to interact with the virtual environment **502**, the distributed processing engine **512** and the data intake and query systems **108.sub.1-108.sub.i**. In particular, the SPL service component **506** may be configured to receive a query from an instance of a programming application operating within the virtual environment **502**, perform analyses on the received query, convert the received query (or an altered version) to a directed acyclic graph (DAG), and provide the DAG to the distributed processing engine **512** for execution.

(224) More specifically, the SPL service component **506** analyzes the query to detect and identify

each operator comprising the query (or particular known fragments of the query comprising a sequence of a plurality of operators). Upon identifying each operator or fragment, the SPL service component **506** determines an input and output schema for each operator or fragment in order to verify that the query is capable of being executed. The SPL service **506** ensures compatibility of the output schema of a first operator with the input schema of a second operator that is directly subsequent the first operator. In some embodiments, the input and output schema of the operators may be defined within a library; thus, the SPL service component **506** may obtain the input and output schema of a plurality of operators by referencing a library corresponding to each operator. Further, when one or more of the operators refers to a machine learning model, the SPL service component **506** may reference metadata of the machine learning operator stored within the operator storage **510** (and specifically, the model storage **511** of FIG. 5B), which includes an indication of the input and output schema of the machine learning operator.

(225) In addition to verifying the compatibility of the schemas of each of the operators comprising the query, the SPL service component **506** provides native support for direct integration of machine learning analytics within a pipelined search query. In current technology, query execution services are unable to support machine learning as an operator within a query. However, the data science platform **504**, and specifically the SPL service component **506**, provides for data scientists to include machine learning algorithms or models as operators directly in a query. In addition, the data science platform **504** may automatically perform query alterations (e.g., to improve or optimize the efficiency of its execution) such as applying projects or filters to one or more operators as a result of analyzing the query in its entirety, discussed in more detail below. As the SPL service component **506** parses a query to detect and identify each operator, the SPL service component **506** identifies operators representing machine learning models or algorithms, where such models or algorithms are pre-stored in a storage medium accessible to the SPL service component **506**, such as the model storage **511**. When an operator is identified as representing a machine learning algorithm or model, the SPL service component **506** retrieves the machine learning algorithm or model from storage. The SPL service component **506** then refers to the machine learning algorithm or model when converting the query to a DAG.

(226) By providing the DAG to the distributed processing engine **512**, the SPL service component **506** decouples the execution of the query from the storage of the data on which the query is to operate (“decoupling compute from storage”). In current technology, query execution services may utilize a single compute node to retrieve the data on which a query is to operate and execute the query, where the single compute node may be collocated with storage of the data, which often results in execution inefficiencies. Such an architecture effectively couples the compute and the storage. The compute nodes in such an architecture are far less computationally efficient than the execution enabled by the data science platform **504** as traditional compute nodes (i) have limited compute and storage resources available, and (ii) serialize read operations included in the query, thereby acting as a bottle neck in the execution of the query. As a result, current query execution services are not able to perform batch machine learning operations and analytics of operators within a query cannot be offloaded to alternative or additional compute nodes. However, the data science platform **504**, and specifically the SPL service component **506**, provides for automated scaling of reads (e.g., retrieval of data) through the utilization of a distributed processing engine to perform the execution of the query. By converting the query (and particularly a pipelined search query) to a DAG, the distributed processing engine may break the DAG into tasks, with a read operation (data retrieval) being broken up into a plurality of tasks where each task may be executed by a separate compute node in a parallel manner with the retrieved data being subsequently stitched together (“parallelized reads”). As a result, serialized reads of data no longer act as a bottle neck during execution.

(227) Additionally, the compute nodes of the distributed processing engine may obtain data from any data storage or stream. For instance and with reference to FIG. 5B, a first compute node, e.g., task

manager **530.sub.1**, of the distributed processing engine **512** may retrieve data from a data store within any of the data intake and query systems **108.sub.1-108.sub.i**, a second compute node, e.g., a second task manager, may retrieve data from any of the data lakes **516.sub.1-516.sub.k** (collective illustrated as “data lake **516**” for purposes of clarity), and a third compute node, e.g., a third task manager, may retrieve data from another location, such as a public cloud storage resource available, which may be AMAZON WEB SERVICES (AWS)® SIMPLE STORAGE SERVICE (S3)® (often referred to as a “S3 bucket”) **517**.

(228) Further, the SPL service component **506** may also perform operations to analyze the query to identify alterations that, if made, would improve the efficiency of the execution of the query. The model service component **508** may be configured to receive a representation of a machine learning model, and specifically, a representation provided in a standardized format, verify the schema of the machine learning model and validate the representation. In some embodiments, the standardized format may be an open-source format, such as the Open Neural Network Exchange (ONNX), which is a format for representing machine learning models through a defined common set of operators and a common file format. As one example, a schema verification process may include deserializing an ONNX model into an in-memory computation graph and comparing the in-memory computation graph to schema metadata, which may be provided by the data scientist or is otherwise accessible by the data science platform **504**. As one particular example, when a provided schema consists of two columns (e.g., both double), then the corresponding in-memory graph should take a vector of size two as input, which would be revealed during a comparison of the schema and the in-memory computation graph. Thus, a verified ONNX model in such an embodiment would have a vector of size two as input (e.g., a match of the schema and the in-memory computation graph). Similarly, when the schema is more complex (e.g., doubles, lists, etc.), the ONNX models will be expected to map to tensors of appropriate shape in the computation graph, where a tensor is an algebraic object that describes a (multilinear) relationship between sets of algebraic objects related to a vector space. And the order of inputs should match the types defined in the schema.

(229) Further, the model service component **508** validates the representation of the machine learning model by attempting to deserialize the representation into one or more computation graphs, such as a standard TensorFlow graph. When the representation is capable of being deserialized, the model service component **508** deems the representation validated.

(230) Referring now to FIG. 5B, a detailed block diagram of the data processing environment **500** of FIG. 5A is shown in accordance with some embodiments. The illustration of FIG. 5B provides a more detailed view of the interactions of the components of the data processing environment **500** as briefly discussed above. In particular, the FIG. 5B illustrates that the virtual environment **502** may include one or more sandboxed environments **518.sub.1-518.sub.L**, with each configured to instantiate and process one or more instances of a programming application, such as an instance of a web-based application, e.g., a JUPYTER® notebook. In particular, each of the sandboxed environments **518.sub.1-518.sub.L** may be referred to operating as a multi-tenant within a single tenant, wherein the multi-tenant connotation refers to the multiple instances of a web-based application operating within a single sandboxed environment **518.sub.1-518.sub.L**. For example, each sandboxed environment **518.sub.1-518.sub.L** may be considered as a customer or client deployment with each of the instances **520.sub.1-520.sub.q** being an instance of a programming application, such as an instance of a JUPYTER® notebook, instantiated by or for the customer or client. Further, the storage **522** may represent isolated storage assigned to each sandboxed environment **518.sub.1-518.sub.L**. In one embodiment, the storage **522** may be AMAZON ELASTIC FILE SYSTEM® (EFS) storage. Further, each of the sandboxed environment **518.sub.1-518.sub.L** may be operating on hardware that is physically separated from the other sandboxed environments. Thus, the customer or client may develop and execute (test) software code within their sandboxed environment.

(231) Additionally, FIG. 5B illustrates a query 524 is received by the SPL service component 506. Upon receipt of the query 524, which may be a pipelined search query, and in some embodiments, a SPL query, the SPL service component 506 parses the query 524 to detect each operator included therein. Each operator is then identified, and the input and output schema of each operator is also identified in order to verify the query 524 will execute based on schema-compatibility among the requisite operators. Following verification of the query 524, the SPL service component 506 may perform operations to alter the query 524 in order to improve the efficiency of its execution. As discussed in more detail below, the SPL service component 506 analyzes the query 524 to determine operators that may be consolidated and represented as a single task (with respect to a DAG, discussed below). Additionally, or in the alternative, alterations to the query 524 may include application of a filter and/or a projection to a data retrieval operator. In addition, alterations may include instructing data retrieval operations to be executed at specific data sources (e.g., the indexing system 112 of FIG. 1). As one example, streaming computations of a single MapReduce sort operation may be executed by the indexing system 112. In addition, some embodiments include improvements that limit the memory usage or reduce CPU use by certain aspects of the query (e.g., of particular algorithms or logic operations included in the query). As one example, when a query is executing percentile calculations on a data set (e.g., 95%, 99%, etc.), alterations may include an instruction to include only a single all-quantile sketch as opposed to creating multiple sketches, and where multiple quantiles are computed from the single sketch. Such an alteration (or alterations) may be performed because such operators are additive and commutative, which permits the combination of such operators when a common grouping key is shared.

(232) Further, the SPL service component 506 performs operations to convert the query 524 to a DAG 525, which is then provided to the distributed processing engine 512, and specifically, a job manager 528.sub.1-528.sub.n (where $n \geq 1$). The job manager 528.sub.1 processes the DAG 525 to provide individual tasks to one or more task managers 530.sub.1-530.sub.p (where $p \geq 1$). The task managers 530.sub.1-530.sub.p execute tasks in accordance with the DAG 525 by, for example, retrieving data as set forth in the query 524 and performing one or more operations on the data. As the results generated by one task manager 530.sub.1-530.sub.p may be provided to another for processing on the results (e.g., as would occur in a pipelined search query).

(233) In some embodiments, the data retrieved from the data intake and query system(s) 108.sub.1-108.sub.i, the data lake 516 or other source may be stored locally by the distributed processing engine 512. For example, when portion of data is routinely retrieved for query execution, that data may be retrieved and cached on storage of or associated with the distributed processing engine 512, where the cache storage may be static random access memory (RAM), also referred to as SRAM 531.

(234) FIG. 5B also illustrates a request 526 for a machine learning model to be stored by the data science platform 504 (at times referred to as “publishing a machine learning model”), and particularly in the model storage 511. The request 526 is received by the model service component 508, where the request may be the result of a call to an API from an instance 520.sub.1-520.sub.q within any of the sandboxed environments 518.sub.1-518.sub.L. The request 524 may include a location at which the model or a representation of the model is stored. The representation may be a language-independent, standardized-format representing the machine learning model, and in some embodiments, the representation may be an ONNX model.

(235) Prior to storage of the representation by the data science platform 504, the model service component 508 performs operations to verify the schema of the model and validate the representation, as discussed above. To validate the representation, the model service component 508 determines whether the representation is capable of being deserialized into a computational graph. When the representation is capable of being deserialized, the representation is deemed validated; otherwise, the representation is not is not deemed validated. Following the verification of the schema and validation of the representation, the model service component 508 stores the

validated representation in the model storage **511**.

(236) Referring to FIG. 5C, an illustration of an operator sequence as a logical plan is shown in accordance with some embodiments. FIG. 5C illustrates a pipeline **532** that includes a plurality of commands such as SPL commands (or operators), where one or more commands may correspond to a request access to a ML component configured to conduct a certain task or tasks, such as a ML model or a ML streaming algorithm for example. The request for access to the ML component may be in the form of a function call to access the binary (executable) code associated with the ML component maintained within the storage system **116** of FIG. 1 (e.g., shared (cloud) data store such as cloud-based storage).

(237) The pipeline **532** tends to start with a first SPL operator **536** associated with a search term or filter criteria. Such search term(s) or filter criteria can include any combination of keywords, phrases, times, dates, Boolean expressions, fieldname-field value pairs, etc. The results can then be passed as input into a second (subsequent) SPL operator **538**, which may include directives for additional processing of the results or commands that assist in the performance of filtering of unwanted information out of the results, extracting information, calculating statistics, reordering the results produced by the first SPL operator **536**, creating an alert or a summary inclusive of a graph, chart, metric, or other visualization of the result data.

(238) As also shown in FIG. 5C, some or all of the pipeline **532** can be separated by a delimiter **534** such as the pipe symbol “|” for example, which may be at least partially relied upon to determine the SPL operators **536**, **538**, **540** and **542**. The parsing of the pipeline **532** results in the detection and identification of operator types (i.e., from command **536**, select command **538**, stats command **540** and detect_drift command **542**), where parsing of the pipeline **532** is detailed in U.S. patent application Ser. No. 17/074,280, now U.S. Pat. No. 11,567,735, entitled “SYSTEMS AND METHODS FOR INTEGRATION OF MULTIPLE PROGRAMMING LANGUAGES WITHIN A PIPELINED SEARCH QUERY”, U.S. patent application Ser. No. 17/074,407, now U.S. Pat. No. 12,131,233, entitled “SYSTEMS AND METHODS FOR AUTO-DEPLOYMENT OF A MACHINE LEARNING COMPONENT WITHIN A PIPELINED SEARCH QUERY” and U.S. patent application Ser. No. 17/704,206, now U.S. Pat. No. 11,748,634, entitled “SYSTEMS AND METHODS FOR INTEGRATION OF MACHINE LEARNING COMPONENTS WITHIN A PIPELINED SEARCH QUERY TO GENERATE A GRAPHIC VISUALIZATION”, each filed concurrently herewith, the contents of each are incorporated by reference herein.

(239) FIG. 5C further illustrates a logical graph **544** corresponding to the pipeline **532**, where each SPL operator is represented by a node with edges between the nodes indicating a particular sequential flow of data between nodes. For example, the logical graph **544** indicates that the result of the read_index operation (SPL operator **546**) is fed into the select operation (e.g., a projection) (SPL operator **548**), the result of which is fed into the stats operation (SPL operator **550**), the result of which is fed into the detect_drift operation (SPL operator **552**).

(240) In current technology, the logical graph **544** is executed as follows: data is read from an index for a given time period into a cluster computing environment, where all data is retrieved from the index regardless of the fields required downstream (i.e., subsequent operators); project the fields timestamp, input and metric (to reduce the retrieved data set); aggregate the relevant data (i.e., streaming quantile by key); and pass the aggregation to the drift_detection operator (e.g., a machine learning algorithm). As noted, one inefficiency of the execution of the logical graph **544** by the current technology is the retrieval (i.e., moving) of all data from the index regardless of the fields required by subsequent operators.

(241) However, the data science platform **504** performs operations to alter the query (e.g., pipeline **532**) that result in an improvement in the efficiency of the execution of the query. For example, data science platform **504** may alter the pipeline **532** to push the projections to the data source (e.g., the indexing system **112**) to only retrieve relevant columns. Additionally, the data science platform **504** may also push the streaming quantile (e.g., the operations performed by the stats operator) to

the data source. In order to perform such alterations, the data science platform **504** analyzes the cost (e.g., CPU or processor usage, memory usage, etc.) of the queries and operations performed by each operator. One advantage of performing such alterations is that, in many situations, only a small fraction of the data at the data source is actually returned to the compute cluster as a result and the compute stage of the graph only needs to scale for certain operations, e.g., machine learning operations such as the drift_detection. The alteration process is detailed in U.S. patent application Ser. No. 17/074,100, now U.S. Pat. No. 11,500,871, entitled “SYSTEMS AND METHODS FOR DECOUPLING SEARCH PROCESSING LANGUAGE AND MACHINE LEARNING ANALYTICS FROM STORAGE OF ACCESSED DATA”, filed concurrently herewith, the contents of which are incorporated by reference herein.

(242) 5.0.2. Operational Flow

(243) Referring to FIGS. **6A-6B**, block diagrams illustrating a plurality of operations performed within the data processing environment **500** of FIG. **5A** are shown in accordance with some embodiments. A first set of operations as illustrated by FIG. **6A** includes operations associated with the validation of a representation of a machine learning model and storage thereof. In particular, user input, such as software code comprising a machine learning model, is received by an instance **520.sub.q** of a web-based application, such as a JUPYTER® notebook (numeral **1**). For purposes of clarity only a single instance of a web-based application is illustrated (the instance **520.sub.q**), but as illustrated in FIG. **5B**, each of the sandboxed environments **518.sub.1-518.sub.L** may have instantiated therein one or more instances. Upon receipt of further user input indicating a request for the machine learning model to be stored by the data science platform **504**, the software code comprising the machine learning model is converted to a representation of the machine learning model being a language-independent, standardized format, such as to an ONNX model. The request for storage of the machine learning, which may include a location at which a representation of the model is stored, is provided to the data science platform **504**, and particularly, to the model service component **508** (numeral **2**). In some embodiments, the API call may result in the conversion of the machine learning model (e.g., a machine learning pipeline) to the ONNX model, where the conversion may or may not be performed by the data science platform **504**.

(244) Upon receipt of the representation of the machine learning model, the model service component **508** performs operations to verify the schema of the machine learning model and further operations to validate the representation (numeral **3**). Following the verification of the schema of the machine learning model and the validation of the representation, the validated representation is stored in the model storage **511** or other storage medium accessible by the model service component **508** and the SPL service component **506** (numeral **4**).

(245) A second set of operations as illustrated by FIG. **6B** includes operations associated with the receipt of a query to be executed, processing of the query, which may include verification, alteration and conversion to a DAG, and providing the DAG to a distributed processing engine. In particular, user input such as software code comprising a query is received by an instance of a web-based application (numeral **5**). Upon receipt of further user input indicating that the query is to be executed, the query is provided to the data science platform **504**, and particularly, to the SPL service component **506** (numeral **6**). The SPL service component **506** performs processing operations including verifying schema-compatibility between operators comprising the query, which may include detecting and identifying each operator and determining the input and output schema for each operator (numeral **7**). In some embodiments, when an operator represents a machine learning model where the representation of which is stored by the data science platform **504**, the SPL service component **506** obtains at least the metadata associated with the representation in order to verify the input and output schema (numeral **8**).

(246) Following verification of the query, the SPL service component **506** further analyzes the query to determine potential alterations to the query that would improve the efficiency of the execution of the query, performs such alterations and converts the query, or altered query, to a DAG

(numeral 9). One such alteration may include application of a filter and/or a projection to a data retrieval operator (e.g., a “read” command) based on data required by the query to reduce the amount of data retrieved. The DAG is then provided to a distributed processing engine 512 (numeral 10). The distributed processing engine 512 may break the DAG into tasks and executes the tasks on a plurality of task managers. In some embodiments, the results of a first task generated by a first task manager are provided to a second task manager such that the results of first task are utilized in performance of the second task.

(247) Referring to FIGS. 7A-7B, flow diagrams illustrating an embodiment of securing a third-party code and processing of a SPL query are shown in accordance with some embodiments. Each block illustrated in FIGS. 7A-7B represents an operation performed in the method 700 performed by a data science platform, such as the data science platform 504 of FIGS. 5A-5B. It should be understood that not every operation illustrated in FIGS. 7A-7B is required. In fact, certain operations may be optional to complete aspects of the method 700. The discussion of the operations of method 700 may be done so with reference to any of the previously described figures. Prior to the initiation of the method 700, it is assumed that one or more instances of a programming application, such as a web-based programming application, have been instantiated and are communicatively coupled to the data science platform 504 of at least FIGS. 5A-5B. In particular, as discussed with respect the embodiment of FIGS. 5A-5B, an instance of a web-based application may be operating within a virtual environment and be configured to transmit a request for storage of a machine learning model as well as a pipelined search query to the data science platform 504.

(248) Referring to FIG. 7A, the method 700 begins with the receipt of user input being software code comprising a machine learning model (block 702). As one example, the software code may be received by an instance of a web-based application, such as a JUPYTER® notebook. The software code may be code comprising a machine learning model or a representation of a machine learning model (e.g., an ONNX model representing the machine learning model). The instance of the web-based application may further receive user input indicating that the machine model is to be published (which may refer to storage of the machine learning model by the data science platform 504) (block 704). As the instance of the web-based application is in communication with the data science platform 504, the data science platform science 504 receives the indication of publication along with a representation of the machine learning model.

(249) Upon receipt of the representation of the machine learning model, the data science platform 504 performs operations to verify the schema of the machine learning model and further operations to validate the machine learning model (block 706).

(250) Following the verification of the schema of the machine learning model and the validation of the representation of the machine learning model, the validated representation is stored by the data science platform 504 (block 708). In some embodiments, the validated representation may be stored on non-transitory storage medium, such as the model storage 511 of FIGS. 5A-5B. In other embodiments, the validated representation may be stored on alternative non-transitory storage medium that is accessible by components of the data science platform 504.

(251) The method 700 continues with the receipt of further software code comprising a pipelined search query, e.g., a SPL query (block 710). As with receipt of the software code comprising a machine learning model (or representation), the instance of the web-based application may receive the SPL query; however, it need not be the same instance that receives both of the machine learning model and SPL query. Referring now to FIG. 7B, further user input, such as additional software code, is received by the web-based application indicating that the SPL query is to be executed (block 712).

(252) The data science platform 504 then receives the SQL query for analysis (block 714). The analysis of the SQL query may include detection and identification of each operator comprising the SQL query and a determination of the schema of each operator. In some embodiments, the schema of the operators may be defined within a library; thus, the data science platform 504 may obtain the

input and output schema of a plurality of operators by referencing a corresponding library of each operator. Further, when one or more of the operators refers to a machine learning model, the data science platform **504** may retrieve the machine learning operator from storage as well as any metadata of the machine learning operator (block **716**). The metadata of a machine learning operator may include the input and output schema of the machine learning operator.

(253) Following the analysis to detect and identify each operator and determining each operator's input and output schema, the data science platform **504** may further analyze the SQL query to determine potential alterations to the SQL query that would improve the efficiency of the execution of the SQL query. As discussed above, one example of alterations to the SQL query may include application of a filter and/or a projection to a data retrieval operator based on data required by the SQL query thereby reducing the amount of data retrieved, which inherently improves the efficiency of execution. The data science platform **504** then alters (e.g., “re-writes”) the SQL query and converts the altered SQL query to a directed acyclic graph (DAG), and provides the DAG to a clustered, distributed processing engine (block **720**). In some embodiments, the clustered, distributed processing engine may be auto-scaling (e.g., automatically laterally scales compute nodes utilized in execution). In some specific embodiments, the clustered, distributed processing engine may be APACHE FLINK®, which is known in the art as a distributed streaming data-flow engine configured to execute dataflow programs in a data-parallel or pipelined manner. The results of the execution of the SQL query may be accessible by the instance of the web-application at a data sink, e.g., a data reservoir, indicated in the SQL query.

(254) 5.0.3. Securement of Untrusted Code

(255) As discussed above, one way that a front-end component such as a JUPYTER® notebook instance **520.sub.1** of FIG. 5B may interact with the data science platform **504** is through a request for a ML model to be stored by the data science platform **504**. As referenced above, the request is received by the model service component **508**, where the request may be the result of a call to an API from a JUPYTER® notebook instance **520.sub.1**. Thus, the receipt of the computation graph **800** (e.g., a ML model representation) by the model service **508** may correspond to the request **526** illustrated in FIG. 5B. Referring now to FIG. 8, a block diagram illustrating an embodiment of the model service component illustrated in FIG. 5A is shown in accordance with some embodiments.

(256) The model service component **508** is deployed as part of the data science platform **504** of FIG. 5A and includes one or more components to receive a ML model representation (e.g., a computation graph) **800**, verify the schema and validate the ML model. In some embodiments, the computation graph **800** may be a known, standard, language-independent format such as the ONNX format (or container). In such embodiments, the conversion of a ML model to a computation graph may be referred to serialization of the ML model to an ONNX model. As is understood, serialization includes converting the ML model (an object) into a stream of bytes so that the ML model may be transferred over a network or stored in non-transitory, computer-readable medium storage. As is known in the art, the ONNX format provides a specific definition of an extensible computation graph model. Importantly, the ONNX format provides definitions for built-in operators and standard data types from which an ONNX model is created. Specifically, an ONNX model is a computation dataflow graph that is structured as a list of nodes forming an acyclic graph. Nodes include one or more inputs and one or more outputs with each node being a call to a function (also referred to as an “operator” by those in the art). An ONNX model also has metadata, which may include inter alia input/output schema of the ML model. Functions may be implemented externally relative to the graph, but the set of built-in functions are specifically designed and configured to be portable across programming language framework. Programming language frameworks that support the ONNX format provide implementations of these functions on the applicable data types.

(257) Additionally, the model service logic **508** may be configured with one or more components to

(i) assist in verifying the input schema of the computation graph **800** matches the input schema of the ML model from which the computation graph was generated and (ii) validate the ML model through deserializing the computation graph **800**. For example, at the time the computation graph **800** was generated from a ML model, the input schema of the ML model is provided at metadata along with the computation graph **800**. Thus, the first step performed by the model service component **508** is to verify that the input schema of the ML model, obtained by accessing the metadata **801**, matches the input schema indicated by the computation graph **800** (in some embodiments, this may be deduced from the training of the ML model).

(258) In the event that the input schema of the ML model as set forth in the metadata **801**, does not match the input schema as indicated by the computation graph **800**, a verification error indicating a schema mismatch is provided to the reporting component **808** such that it may be provided to, for example, a data scientist, that developed the ML pipeline and model. The schema verification error may be provided as any form of a message, such as via a visual alert presented by a user interface of a network device. This allows the data scientist to identify the schema mismatch and correct prior to resubmission.

(259) Following a successful schema verification in which the input schemas match, the model validation component **804** performs at least a deserialization of the computation graph **800** and upon completing the deserialization without obtaining one or more errors, the computation **800** is stored in model **511**. Examples of deserialization errors may include, but are not limited or restricted to, a failure to provide definitions, missing directive or assembly reference, mismatched argument types, etc. In some embodiments, following a successful deserialization, the deserialized ML model object may be reserialized to confirm that a resulting reserialized computation graph matches the received computation graph **800**. Of course, when either the deserialization of the computation graph **800** or the reserialization fails, a validation error may be provided to the reporting component **808** such that it may be provided to, for example, a data scientist, that developed the ML pipeline and model.

(260) Referring now to FIG. **9**, a block diagram illustrating an embodiment of the components associated with the model validation component **804** deployed within the model service component **508** illustrated in FIG. **8** is shown in accordance with some embodiments. Herein, the model validation component **804** features a deserialization component **900**, and an optional serialization component **902**.

(261) Following verification of the model schema by the model schema verification component **802**, the computation graph **800** is provided to the model validation component **804**, and particularly, the deserialization component **900**, which attempts to perform a deserialization of the computation graph **800**. As is understood, deserialization is the reverse process of serialization and includes converting the computation graph **800** (a stream of bytes) into a ML model (an object representation of the ML model) **904**. By deserializing the computation graph, the model validation component **804** may confirm whether the deserialized ML model is a legitimate and valid ML model. Specifically, as the computation graph **800** is comprised of only a predefined set of built-in functions (or operators), the object representation of a ML model obtained from deserializing the computation graph **800** is considered legitimate and valid. Specifically, malicious or anomalous code cannot be inserted into a computation graph due to the limited set of built-in functions; therefore, the successful deserialization of a computation graph can be used as an indicator that the computation graph is valid. Thus, by successfully deserializing the computation graph into a ML model, the model validation component **804** validates the ML model and the corresponding computation graph, and stores the original serialized computation graph. In some embodiments, the deserialization component **900** may store the original, serialized computation graph **800** in the model storage **511**.

(262) As one non-limiting example, the deserialization by the deserialization component **900** results in a JAVA® object representation of the ML model. Deserialization into a JAVA® object

representation may be utilized by the data science platform **504** for processing, e.g., altering a received query that calls the ML model in improve the efficiency of the execution of the query, as described in U.S. patent application Ser. No. 17/704,441, entitled “SYSTEMS AND METHODS FOR SECURING UNTRUSTED CODE” filed concurrently herewith, the entire contents of which are incorporated by reference herein.

(263) However, when the deserialization fails, the deserialization issue may optionally be identified and a validation error may be generated and provided to the reporting component **808**. The validation error may in turn be provided to, for example, a data scientist, that developed the ML pipeline and model. In some embodiments, the validation error may be provided in any form such as that of a message, e.g., via a visual alert presented by a user interface of a network device.

(264) In some embodiments, which may be alternative to the automatic storage of the computation graph **800** following a successful deserialization, the deserialized ML model **904** may be provided to the serialization component **902**, which is configured serialize (reserialize) the deserialized ML model **904** (object representation) into a computation graph **905**. The serialization component **902** then compares the original serialized computation graph **800** to the reserialized computation graph **905** to determine whether the two match. A match would serve to further indicate that malicious code (or anomalous or otherwise unwanted code) was not inserted into the ML model or the original computation graph **800**. Thus, in the event of a match, the original serialized computation graph **800** is stored on or by the data science platform in, for example, the model storage **511**.

(265) However, in the event that the reserialization was unsuccessful or the comparison the of original serialized computation graph **800** with the reserialized computation graph **905** did not yield a match, a validation error may be generated and provided to, for example, a data scientist, that developed the ML pipeline and model via the reporting component **808**. The validation error may be provided in the same manner as discussed above.

(266) Referring to FIGS. **10A-10C**, flow diagrams illustrating an embodiment of securing a machine-learning model for storage on the data science platform of FIG. **5A** are shown in accordance with some embodiments. Each block illustrated in FIGS. **10A-10C** represents an operation performed in the method **1000**, where a plurality of the operations are performed by a data science platform, such as the data science platform **504** of FIGS. **5A-5B**. It should be understood that not every operation illustrated in FIGS. **10A-10C** is required. In fact, certain operations may be optional to complete aspects of the method **1000**. Specifically, the entire set of operations of FIG. **10C** may be optional. The discussion of the operations of method **1000** may be done so with reference to any of the previously described figures. Prior to the initiation of the method **1000**, it is assumed that one or more instances of a programming application, such as a web-based programming application, have been instantiated and are communicatively coupled to the data science platform **504** of at least FIGS. **5A-5B**. In particular, as discussed with respect the embodiment of FIGS. **5A-5B**, an instance of a web-based application may be operating within a virtual environment and be configured to transmit a request for storage of a machine learning model (or representation thereof) to the data science platform **504**.

(267) Referring to FIG. **10A**, the method **1000** begins with the development of a machine learning (ML) pipeline using a plurality of operators, where development of the ML pipeline may be a result of receipt of user input being software code (block **1002**). Following development of the ML pipeline, a ML model is generated through training with the ML pipeline and the ML model is converted to a computation graph (block **1004**). The conversion of the ML model to a computation graph may be referred to as serializing the ML model. The computation graph may be a known, standard, language-independent format such as the ONNX format (or container). In such embodiments, the conversion of a ML model to a computation graph may be referred to as conversion to an ONNX model. As is understood, serialization includes converting the ML model (an object) into a stream of bytes so that the ML model may be transferred over a network or stored in non-transitory, computer-readable medium storage.

(268) As is known in the art, the ONNX format provides a specific definition of an extensible computation graph model. Importantly, the ONNX format provides definitions for built-in operators and standard data types from which an ONNX model is created. Specifically, an ONNX model is a computation dataflow graph that is structured as a list of nodes forming an acyclic graph. Nodes include one or more inputs and one or more outputs with each node being a call to a function (also referred to as an “operator” by those in the art). An ONNX model also has metadata, which may include inter alia input/output schema of the ML model. Functions may be implemented externally relative to the graph, but the set of built-in functions are specifically designed and configured to be portable across programming language framework. Programming language frameworks that support the ONNX format provides implementations of these functions on the applicable data types.

(269) The method **1000** continues as the data science platform **504** receives the computation graph and verifies the input schema of the computation graph matches (or is otherwise compatible) the input schema deduced from the training of the ML model, which may be part of the metadata provided along with the computation graph (such as the metadata **801** of FIG. **8**) (blocks **1006** and **1008**).

(270) When the input schemas do not match, the mismatch between the input schemas may optionally be identified and a computation graph schema verification error may be generated and provided to, for example, a data scientist, that developed the ML pipeline and model (blocks **1010** and **1012**). The computation graph schema validation error may be provided in any form of a message, e.g., via a visual alert presented by a user interface of a network device. This allows the data scientist to identify the schema mismatch and correct prior to resubmission. Further, this may indicate an instability or vulnerability within the ML model or computation graph. Alternatively, or in addition, this may indicate that malicious code has been inserted into the ML model and/or the computation graph; thus, such computation graphs are not stored on or by the data science platform **504**. Specifically, prevention of storage of ML models with schema mismatches from being stored on or by the data science platform **504** is one step in securing the ML model and ensuring malicious or otherwise unstable code is not executed by the data science platform **504**.

(271) Referring to FIG. **10B**, the method **1000** continues when the input schemas match with a deserialization of the computation graph into an object representation of the ML model (block **1014**). As is understood, deserialization is the reverse process of serialization and includes converting the computation graph (a stream of bytes) into a ML model (an object representation of the ML model). By deserializing the computation graph, the model validation component **804** may confirm whether the deserialized ML model is a legitimate and valid ML model. Specifically, as a computation graph is comprised of only a predefined set of built-in functions (or operators), a ML model obtained from deserializing a computation graph is considered legitimate and valid as malicious or anomalous code cannot be inserted into a computation graph. Thus, by successfully deserializing the computation graph into a ML model, the model validation component **804** validates the ML model and the corresponding computation graph, and stores the original serialized computation graph (blocks **1016** and **1022**).

(272) However, when the deserialization fails, the deserialization issue may optionally be identified and a validation error may be generated and provided to, for example, a data scientist, that developed the ML pipeline and model (blocks **1016**, **1018** and **1020**). The validation error may be provided in the same manner as discussed above with respect to the computation graph schema validation error.

(273) In addition to the operations illustrated in FIGS. **10A-10B**, the operations of FIG. **10C** illustrate an alternative, or additional, embodiment. As illustrated in FIG. **10B**, following a determination that deserialization was unsuccessful (block **1016**), a first embodiment of the method **1000** may include storing the original serialized computation graph (block **1022**). FIG. **10B** also indicates that an alternative, second embodiment is illustrated in FIG. **10C**, where, following

successful deserialization of the computation graph into an object representation of the ML model, the object representation is reserialized into a computation graph (block **1024**). In one particular embodiment, the deserialization results in a JAVA® object representation of the ML model, where the data science platform **504** utilizes the JAVA® object representation for processing, e.g., when the data science platform **504** receives a query including a call to the ML model, such as the query **524** of FIG. **5B**. Thus, in such embodiments, deserializing a received computation graph serves as a first step in validating the computation graph (and ML model).

(274) Additionally, the operations of reserializing the JAVA® object representation into a computation graph, e.g., an ONNX model, and comparing the original serialized computation graph to the reserialized computation graph serves as a second step in validating the computation graph (and ML model) (blocks **1026** and **1030**). When the original serialized computation graph and the reserialized computation graph match, the original serialized computation graph may be stored on or by the data science platform (blocks **1034** and **1036**).

(275) However, when the reserialization was unsuccessful or the comparison of the original serialized computation graph with the reserialized computation graph did not yield a match, a validation error may be generated and provided to, for example, a data scientist, that developed the ML pipeline and model (blocks **1026**, **1028**, **1030**, **1032** and **1034**). The validation error may be provided in the same manner as discussed above.

(276) 6.0. Terminology

(277) Computer programs typically comprise one or more instructions set at various times in various memory devices of a computing device, which, when read and executed by at least one processor, will cause a computing device to execute functions involving the disclosed techniques. In some embodiments, a carrier containing the aforementioned computer program product is provided. The carrier is one of an electronic signal, an optical signal, a radio signal, or a non-transitory computer-readable storage medium.

(278) Any or all of the features and functions described above can be combined with each other, except to the extent it may be otherwise stated above or to the extent that any such embodiments may be incompatible by virtue of their function or structure, as will be apparent to persons of ordinary skill in the art. Unless contrary to physical possibility, it is envisioned that (i) the methods/steps described herein may be performed in any sequence and/or in any combination, and (ii) the components of respective embodiments may be combined in any manner.

(279) Although the subject matter has been described in language specific to structural features and/or acts, it is to be understood that the subject matter defined in the appended claims is not necessarily limited to the specific features or acts described above. Rather, the specific features and acts described above are disclosed as examples of implementing the claims, and other equivalent features and acts are intended to be within the scope of the claims.

(280) Conditional language, such as, among others, “can,” “could,” “might,” or “may,” unless specifically stated otherwise, or otherwise understood within the context as used, is generally intended to convey that certain embodiments include, while other embodiments do not include, certain features, elements and/or steps. Thus, such conditional language is not generally intended to imply that features, elements and/or steps are in any way required for one or more embodiments or that one or more embodiments necessarily include logic for deciding, with or without user input or prompting, whether these features, elements and/or steps are included or are to be performed in any particular embodiment. Furthermore, use of “e.g.,” is to be interpreted as providing a non-limiting example and does not imply that two things are identical or necessarily equate to each other.

(281) Unless the context clearly requires otherwise, throughout the description and the claims, the words “comprise,” “comprising,” and the like are to be construed in an inclusive sense, as opposed to an exclusive or exhaustive sense, i.e., in the sense of “including, but not limited to.” As used herein, the terms “connected,” “coupled,” or any variant thereof means any connection or coupling, either direct or indirect, between two or more elements; the coupling or connection between the

elements can be physical, logical, or a combination thereof. Additionally, the words “herein,” “above,” “below,” and words of similar import, when used in this application, refer to this application as a whole and not to any particular portions of this application. Where the context permits, words using the singular or plural number may also include the plural or singular number respectively. The word “or” in reference to a list of two or more items, covers all of the following interpretations of the word: any one of the items in the list, all of the items in the list, and any combination of the items in the list. Likewise the term “and/or” in reference to a list of two or more items, covers all of the following interpretations of the word: any one of the items in the list, all of the items in the list, and any combination of the items in the list.

(282) Conjunctive language such as the phrase “at least one of X, Y and Z,” unless specifically stated otherwise, is understood with the context as used in general to convey that an item, term, etc. may be either X, Y or Z, or any combination thereof. Thus, such conjunctive language is not generally intended to imply that certain embodiments require at least one of X, at least one of Y and at least one of Z to each be present. Further, use of the phrase “at least one of X, Y or Z” as used in general is to convey that an item, term, etc. may be either X, Y or Z, or any combination thereof.

(283) In some embodiments, certain operations, acts, events, or functions of any of the algorithms described herein can be performed in a different sequence, can be added, merged, or left out altogether (e.g., not all are necessary for the practice of the algorithms). In certain embodiments, operations, acts, functions, or events can be performed concurrently, e.g., through multi-threaded processing, interrupt processing, or multiple processors or processor cores or on other parallel architectures, rather than sequentially.

(284) Systems and modules described herein may comprise software, firmware, hardware, or any combination(s) of software, firmware, or hardware suitable for the purposes described. Software and other modules may reside and execute on servers, workstations, personal computers, computerized tablets, PDAs, and other computing devices suitable for the purposes described herein. Software and other modules may be accessible via local computer memory, via a network, via a browser, or via other means suitable for the purposes described herein. Data structures described herein may comprise computer files, variables, programming arrays, programming structures, or any electronic information storage schemes or methods, or any combinations thereof, suitable for the purposes described herein. User interface elements described herein may comprise elements from graphical user interfaces, interactive voice response, command line interfaces, and other suitable interfaces.

(285) Further, processing of the various components of the illustrated systems can be distributed across multiple machines, networks, and other computing resources. Two or more components of a system can be combined into fewer components. Various components of the illustrated systems can be implemented in one or more virtual machines or an isolated execution environment, rather than in dedicated computer hardware systems and/or computing devices. Likewise, the data repositories shown can represent physical and/or logical data storage, including, e.g., storage area networks or other distributed storage systems. Moreover, in some embodiments the connections between the components shown represent possible paths of data flow, rather than actual connections between hardware. While some examples of possible connections are shown, any of the subset of the components shown can communicate with any other subset of components in various implementations.

(286) Embodiments are also described above with reference to flow chart illustrations and/or block diagrams of methods, apparatus (systems) and computer program products. Each block of the flow chart illustrations and/or block diagrams, and combinations of blocks in the flow chart illustrations and/or block diagrams, may be implemented by computer program instructions. Such instructions may be provided to a processor of a general purpose computer, special purpose computer, specially-equipped computer (e.g., comprising a high-performance database server, a graphics

subsystem, etc.) or other programmable data processing apparatus to produce a machine, such that the instructions, which execute via the processor(s) of the computer or other programmable data processing apparatus, create means for implementing the acts specified in the flow chart and/or block diagram block or blocks. These computer program instructions may also be stored in a non-transitory computer-readable medium (memory) that can direct a computer or other programmable data processing apparatus to operate in a particular manner, such that the instructions stored in the computer-readable memory produce an article of manufacture including instruction means which implement the acts specified in the flow chart and/or block diagram block or blocks. The computer program instructions may also be loaded to a computing device or other programmable data processing apparatus to cause operations to be performed on the computing device or other programmable apparatus to produce a computer implemented process such that the instructions which execute on the computing device or other programmable apparatus provide steps for implementing the acts specified in the flow chart and/or block diagram block or blocks.

(287) Any patents and applications and other references noted above, including any that may be listed in accompanying filing papers, are incorporated herein by reference. Aspects of the invention can be modified, if necessary, to employ the systems, functions, and concepts of the various references described above to provide yet further implementations of the invention. These and other changes can be made to the invention in light of the above Detailed Description. While the above description describes certain examples of the invention, and describes the best mode contemplated, no matter how detailed the above appears in text, the invention can be practiced in many ways. Details of the system may vary considerably in its specific implementation, while still being encompassed by the invention disclosed herein. As noted above, particular terminology used when describing certain features or aspects of the invention should not be taken to imply that the terminology is being redefined herein to be restricted to any specific characteristics, features, or aspects of the invention with which that terminology is associated. In general, the terms used in the following claims should not be construed to limit the invention to the specific examples disclosed in the specification, unless the above Detailed Description section explicitly defines such terms. Accordingly, the actual scope of the invention encompasses not only the disclosed examples, but also all equivalent ways of practicing or implementing the invention under the claims.

(288) To reduce the number of claims, certain aspects of the invention are presented below in certain claim forms, but the applicant contemplates other aspects of the invention in any number of claim forms. For example, while only one aspect of the invention is recited as a means-plus-function claim under 35 U.S.C. § 112(f) (AIA), other aspects may likewise be embodied as a means-plus-function claim, or in other forms, such as being embodied in a computer-readable medium. Any claims intended to be treated under 35 U.S.C. § 112(f) will begin with the words “means for,” but use of the term “for” in any other context is not intended to invoke treatment under 35 U.S.C. § 112(f). Accordingly, the applicant reserves the right to pursue additional claims after filing this application, in either this application or in a continuing application.

Claims

1. A computer-implemented method, comprising: receiving, by a data processing environment, a request to store a representation of a machine learning (ML) model, wherein the representation of the ML model is a first computation graph, and wherein metadata of the ML model is received with the request, wherein the metadata includes an input schema and an output schema of the first computation graph; performing, by a data processing environment, a verification process on the ML model that includes (i) attempting to deserialize the representation of the ML model by converting the first computation graph into a second computation graph, and (ii) comparing an input schema of the second computation graph with an input schema provided in the metadata, wherein a successful verification includes (1) successfully deserializing the representation of the

ML, and (2) determining that the input schema of the second computation graph matches the input schema provided in the metadata; performing, by a data processing environment, a validation process on the ML model that includes attempting to deserialize the representation of the ML model by converting the first computation graph into a second computation graph, wherein successfully validating the ML model includes completion of the validation process without obtaining a deserialization error; storing, by a data processing environment, the representation of the ML model upon successful completion of the verification process and the validation process; executing, by a data processing environment, a pipelined search query that comprises a sequence of operators including a first operator and a second operator, wherein executing the pipelined search query comprises: includes-applying the first operator to a set of data indicated by the pipelined search query to produce first results; and applying the second operator to the first results to generate second results wherein the ML model is deployed as either the first operator or the second operator; and causing display, by a data processing environment, results of the execution of the pipelined search query on a user interface of a network device.

2. The method of claim 1, wherein the representation of the ML model is a format for representing ML models through a defined common set of operators and a common file format.

3. The method of claim 1, further comprising: identifying a mismatch or an incompatibility between the input schema of the representation of the ML model and the input schema of the ML model.

4. The method of claim 3, further comprising: generating an error report indicating the mismatch or the incompatibility.

5. The method of claim 1, further comprising: generating an error report indicating an error with the deserializing of the representation of the ML model.

6. The method of claim 1, wherein validating the representation of the ML model further includes: serializing the deserialized representation of the ML model, and comparing the deserialized representation of the ML model to the representation of ML model originally received in the request, wherein a match indicates that the representation of the ML model is valid and secure.

7. The method of claim 6, further comprising: generating an error report indicating an error with the serializing of the deserialized representation of the ML model.

8. The method of claim 1, wherein the request to store the representation of the ML model is received from an instance of a web-based application operating within a sandboxed environment.

9. The method of claim 8, wherein the sandboxed environment is configured to have operated therein a plurality of instances of the web-based application, and wherein the sandboxed environment operates with isolated storage assigned thereto.

10. A computing device, comprising: a processor; and a non-transitory computer-readable medium having stored thereon instructions that, when executed by the processor, cause the processor to perform operations including: receiving, by a data processing environment, a request to store a representation of a machine learning (ML) model, wherein the representation of the ML model is a first computation graph, and wherein metadata of the ML model is received with the request, wherein the metadata includes an input schema and an output schema of the first computation graph; performing, by a data processing environment, a verification process on the ML model that includes (i) attempting to deserialize the representation of the ML model by converting the first computation graph into a second computation graph, and (ii) comparing an input schema of the second computation graph with an input schema provided in the metadata, wherein a successful verification includes (1) successfully deserializing the representation of the ML, and (2) determining that the input schema of the second computation graph matches the input schema provided in the metadata; performing, by a data processing environment, a validation process on the ML model that includes attempting to deserialize the representation of the ML model by converting the first computation graph into a second computation graph, wherein successfully validating the ML model includes completion of the validation process without obtaining a

deserialization error; storing, by a data processing environment, the representation of the ML model upon successful completion of the verification process and the validation process; executing, by a data processing environment, a pipelined search query that comprises a sequence of operators including a first operator and a second operator, wherein executing the pipelined search query comprises: applying the first operator to a set of data indicated by the pipelined search query to produce first results; and applying the second operator to the first results to generate second results, wherein the ML model is deployed as either the first operator or the second operator; and causing display, by a data processing environment, results of the execution of the pipelined search query on a user interface of a network device.

11. The computing device of claim 10, wherein the representation of the ML model is a format for representing ML models through a defined common set of operators and a common file format.

12. The computing device of claim 10, further comprising: identifying a mismatch or an incompatibility between the input schema of the representation of the ML model and the input schema of the ML model.

13. The computing device of claim 12, further comprising: generating an error report indicating the mismatch or the incompatibility.

14. The computing device of claim 10, further comprising: generating an error report indicating an error with the deserializing of the representation of the ML model.

15. The computing device of claim 10, wherein validating the representation of the ML model further includes: serializing the deserialized representation of the ML model, and comparing the deserialized representation of the ML model to the representation of ML model originally received in the request, wherein a match indicates that the representation of the ML model is valid and secure.

16. The computing device of claim 10, wherein the request to store the representation of the ML model is received from an instance of a web-based application operating within a sandboxed environment.

17. The computing device of claim 16, wherein the sandboxed environment is configured to have operated therein a plurality of instances of the web-based application, and wherein the sandboxed environment operates with isolated storage assigned thereto.

18. A non-transitory computer-readable medium having stored thereon instructions that, when executed by one or more processors, cause the one or more processor to perform operations including: receiving, by a data processing environment, a request to store a representation of a machine learning (ML) model, wherein the representation of the ML model is a first computation graph, and wherein metadata of the ML model is received with the request, wherein the metadata includes an input schema and an output schema of the first computation graph; performing, by a data processing environment, a verification process on the ML model that includes (i) attempting to deserialize the representation of the ML model by converting the first computation graph into a second computation graph, and (ii) comparing an input schema of the second computation graph with an input schema provided in the metadata, wherein a successful verification includes (1) successfully deserializing the representation of the ML, and (2) determining that the input schema of the second computation graph matches the input schema provided in the metadata; performing, by a data processing environment, a validation process on the ML model that includes attempting to deserialize the representation of the ML model by converting the first computation graph into a second computation graph, wherein successfully validating the ML model includes completion of the validation process without obtaining a deserialization error; storing, by a data processing environment, the representation of the ML model upon successful completion of the verification process and the validation process; executing, by a data processing environment, a pipelined search query that comprises a sequence of operators including a first operator and a second operator, wherein executing the pipelined search query comprises: applying the first operator to a set of data indicated by the pipelined search query to produce first results; and applying the second operator to

the first results to generate second results wherein the ML model is deployed as either the first operator or the second operator; and causing display, by a data processing environment, results of the execution of the pipelined search query on a user interface of a network device.

19. The non-transitory computer-readable medium of claim 18, wherein the request to store the representation of the ML model is received from an instance of a web-based application operating within a sandboxed environment.

20. The non-transitory computer-readable medium of claim 19, wherein the sandboxed environment is configured to have operated therein a plurality of instances of the web-based application, and wherein the sandboxed environment operates with isolated storage assigned thereto.
