



US012394014B2

(12) **United States Patent**
Mehendale et al.

(10) **Patent No.:** **US 12,394,014 B2**

(45) **Date of Patent:** **Aug. 19, 2025**

(54) **APPARATUS AND METHOD FOR
EFFICIENT MOTION ESTIMATION**

(71) Applicant: **TEXAS INSTRUMENTS
INCORPORATED**, Dallas, TX (US)

(72) Inventors: **Mahesh M. Mehendale**, Bangalore
(IN); **Ajit Deepak Gupte**, Bangalore
(IN)

(73) Assignee: **TEXAS INSTRUMENTS
INCORPORATED**, Dallas, TX (US)

(*) Notice: Subject to any disclaimer, the term of this
patent is extended or adjusted under 35
U.S.C. 154(b) by 0 days.

(21) Appl. No.: **18/379,249**

(22) Filed: **Oct. 12, 2023**

(65) **Prior Publication Data**

US 2024/0037700 A1 Feb. 1, 2024

Related U.S. Application Data

(63) Continuation of application No. 17/411,844, filed on
Aug. 25, 2021, now Pat. No. 11,790,485, which is a
(Continued)

(30) **Foreign Application Priority Data**

May 4, 2016 (IN) 201641015446

(51) **Int. Cl.**
G06T 7/238 (2017.01)
G06T 3/4007 (2024.01)

(Continued)

(52) **U.S. Cl.**
CPC **G06T 3/4007** (2013.01); **G06T 7/238**
(2017.01); **H04N 19/43** (2014.11); **H04N**
19/523 (2014.11);

(Continued)

(58) **Field of Classification Search**

CPC G06T 3/4007; G06T 7/238; G06T
2207/10016; G06T 2207/20021; H04N
19/43; H04N 19/523

See application file for complete search history.

(56) **References Cited**

U.S. PATENT DOCUMENTS

6,026,217 A * 2/2000 Adiletta H04N 19/159
709/247
7,782,951 B2 * 8/2010 Muthukrishnan H04N 19/53
375/240.15

(Continued)

FOREIGN PATENT DOCUMENTS

WO WO2016160608 A1 10/2016

OTHER PUBLICATIONS

Grzegorz Pastuszak et al., "Optimization of the Adaptive Compu-
tationally-Scalable Motion Estimation and Compensation for Hard-
ware H.264/AVC Encoder" DOI 10.1007/s 11265-015-1021-5; J
Sign Process Syst. Jul. 21, 2015 at Springerlink.com (Year: 2015)
(Year: 2015).*

(Continued)


Primary Examiner — Dramos Kalapodas

(74) *Attorney, Agent, or Firm* — Carl G. Peterson; Frank
D. Cimino

(57) **ABSTRACT**

The architecture shown can perform global search, local
search and local sub pixel search in a parallel or in a
pipelined mode. All operations are in a streaming mode
without the requirement of external intermediate data stor-
age.

20 Claims, 5 Drawing Sheets



Related U.S. Application Data

continuation of application No. 16/809,052, filed on Mar. 4, 2020, now Pat. No. 11,127,114, which is a continuation of application No. 15/586,600, filed on May 4, 2017, now Pat. No. 10,593,015.

2016/0021385	A1 *	1/2016	Chou	H04N 19/513 375/240.16
2016/0105682	A1 *	4/2016	Rapaka	H04N 19/50 375/240.12
2018/0255311	A1	9/2018	Esenlik et al.	

(51) **Int. Cl.***H04N 19/43* (2014.01)*H04N 19/523* (2014.01)(52) **U.S. Cl.**

CPC *G06T 2207/10016* (2013.01); *G06T 2207/20021* (2013.01)

(56)

References Cited

U.S. PATENT DOCUMENTS

9,699,456	B2	7/2017	Chien et al.	
9,872,044	B2 *	1/2018	Nandan	H04N 19/86
9,918,102	B1	3/2018	Kohn et al.	
10,187,655	B2 *	1/2019	Chou	H04N 19/147
2008/0059467	A1 *	3/2008	Bivolarski	G06F 9/3885 712/E9.05
2013/0208796	A1	8/2013	Amitay et al.	
2015/0084970	A1 *	3/2015	Schaub	H04N 19/423 345/506
2015/0195524	A1 *	7/2015	Li	H04N 19/46 375/240.02

OTHER PUBLICATIONS

Kashi Viswanatha Reddy, "Fast Block Matching Motion Estimation Algorithms for Video Compression", 2013; NIT ROURKELA-769008 (Year: 2013) (Year: 2013).*

Kashi Viswanatha Reddy, "Fast Block Matching Motion Estimation Algorithms for Video Compression", 2013; NIT ROURKELA-769008 (Year: 2013).

Martin Schwalb et al., "Fast Motion Estimation on Graphics Hardware for H.264 Video Encoding", 10.1109/TMM.2008.2008873(c) 2009 IEEE (Year: 2009).

Francis Kelly et al., "Fast Image Interpolation for Motion estimation Using Graphics Hardware", Proc. SPIE 5297, Real-Time Imaging VIII (May 18, 2004); doi: 10.1117/12.526400 (Year 2004).

Lawrence Chan et al., "Parallelizing H.264 Motion Estimation Algorithm using CUDA" IAP-MIT 2009 (Year: 2009).

Grzegorz Pastuszek et al., "Optimization of the Adaptive Computationally-Scalable Motion Estimation and Compensation for Hardware H.264/AVC Encoder" DOI 10.1007/s 11265-015-1021-5; J Sign Process Syst. Jul. 21, 2015 at Springerlink.com (Year: 2015).
Cor Meenderinck et al., "Parallel Scalability of Video DEcoders"; DOI 10.1007/s 11265-008-0256-9 (c) Meendenrinck (Year: 2008).

* cited by examiner

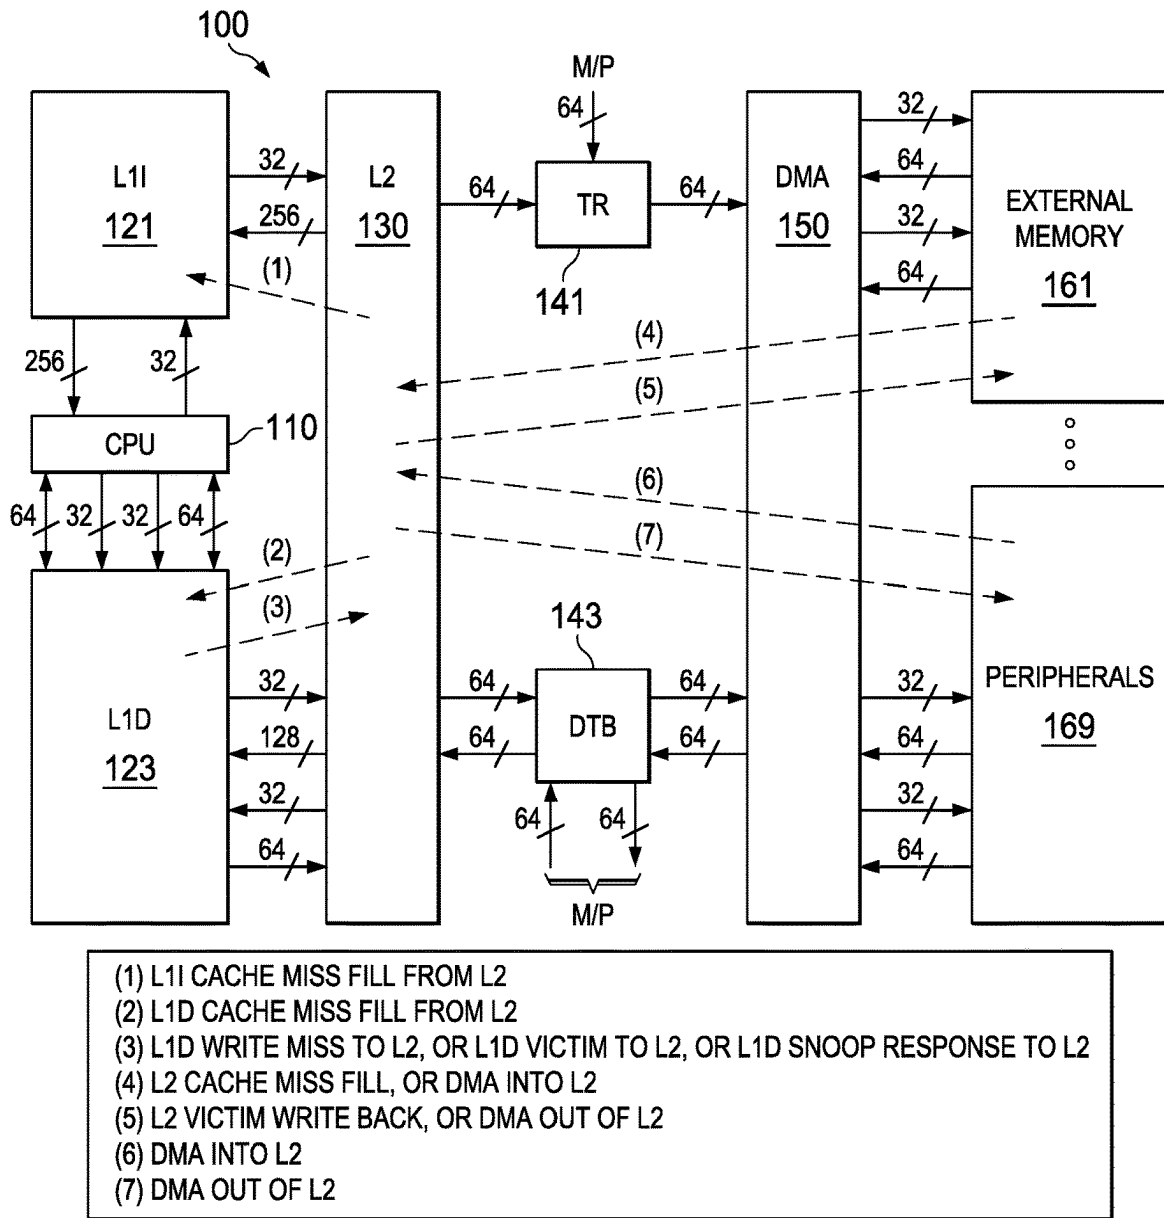



FIG. 1
(PRIOR ART)






FIG. 2
(PRIOR ART)




FIG. 4
(PRIOR ART)

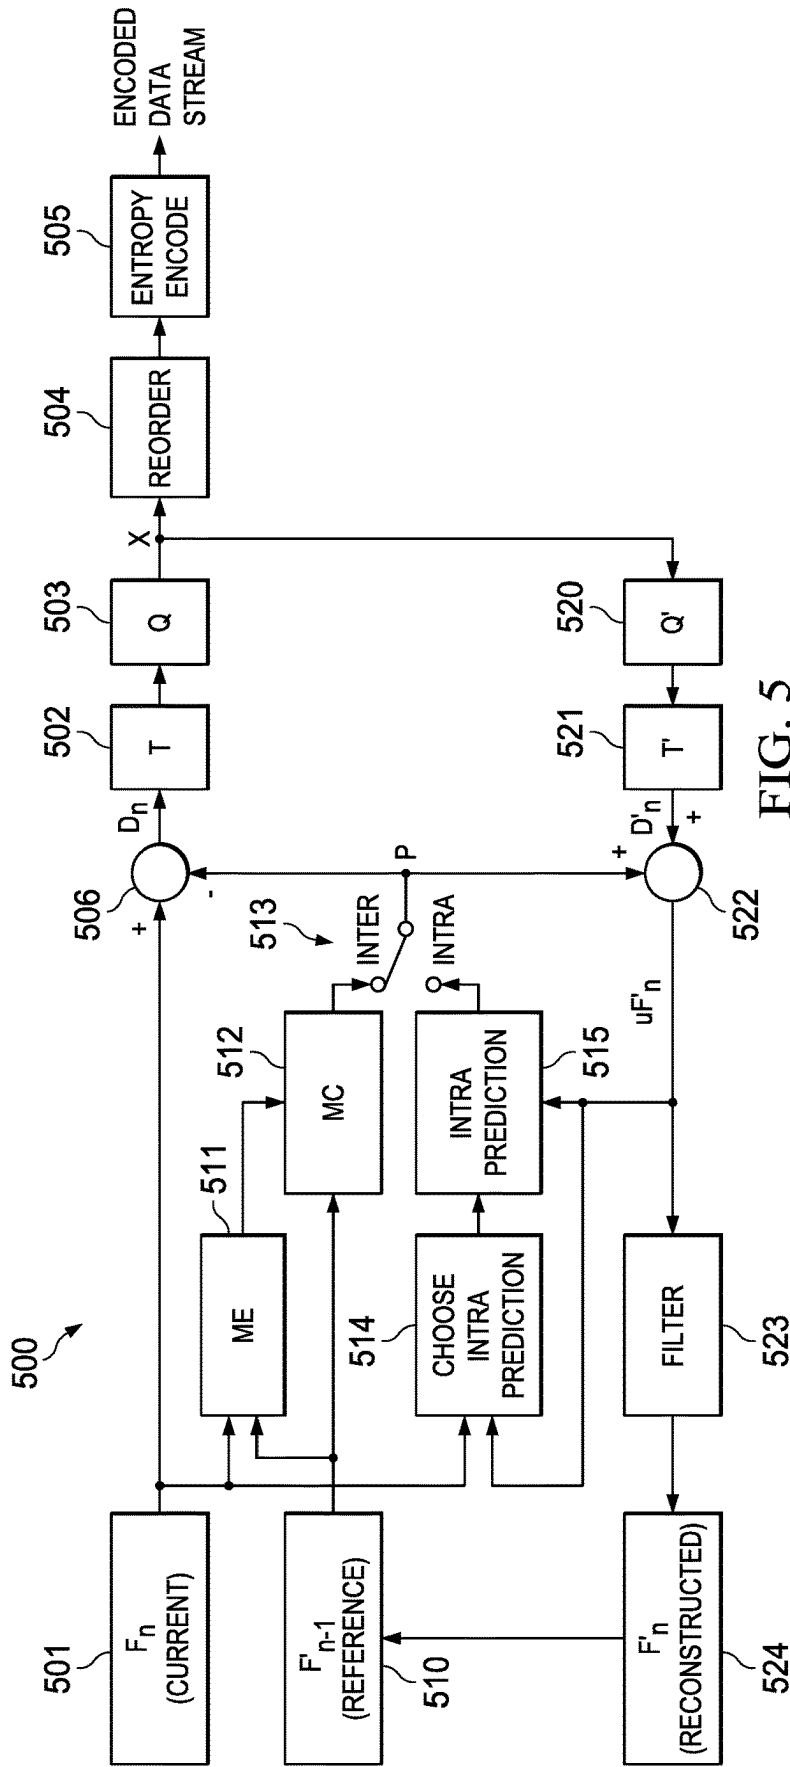





FIG. 5
(PRIOR ART)





APPARATUS AND METHOD FOR EFFICIENT MOTION ESTIMATION

CROSS-REFERENCE TO RELATED APPLICATIONS

This application is a continuation of U.S. patent application Ser. No. 17/411,844, filed Aug. 25, 2021, and scheduled to grant as U.S. Pat. No. 11,790,485 on Oct. 17, 2023, which is a continuation of U.S. patent application Ser. No. 16/809,052, filed Mar. 4, 2020 (now U.S. Pat. No. 11,127,114), which is a continuation of U.S. patent application Ser. No. 15/586,600, filed May 4, 2017 (now U.S. Pat. No. 10,593,015), which claims priority under 35 U.S.C. 119(e)(1) to Indian Provisional Application No. 201641015446, filed May 4, 2016, each of which is herein incorporated by reference in its entirety.

TECHNICAL FIELD OF THE INVENTION

The technical field of this invention is image compression.

BACKGROUND OF THE INVENTION

Increasing video resolution and frame rates, along with large number of searching and matching operations involved in motion estimation demand very high performance. While high performance can be achieved by increasing hardware throughput and higher clock frequency, it is important to identify and exploit parallelism present in the algorithm in order to efficiently utilize available hardware resources.

The Motion Estimation process involves searching operations which require accessing large amounts of reference picture data from memory. Memory bandwidth is an expensive resource which often limits the computational parallelism that can be built in hardware. Further, this large data traffic from the memory leads to large power dissipation.

Motion estimation finds a best match for each block in a current video frame among blocks from previously coded frame(s) (called as reference frames). Block size is typically 16×16 pixels.

A widely used metric to define the match is SAD (Sum Of Absolute Difference in all the pixel values of current block and a reference block).

The best match information is indicated by the motion vector: if the current position of a block is (16,16) then motion vector (4,1) means the best match lies at position (20,17) in the reference frame.

The motion vector can also be in fraction pixel precision: half pixel, quarter pixel etc.

Fractional pixels are calculated by interpolating neighboring integer position pixels.

A motion estimation algorithm would typically include these steps:

- Stage 1: choosing best among a few predictor motion vectors;
- Stage 2: search around winner of Stage 1;
- Stage 3, 4: search around winner of Stage 2 and Stage 3 respectively;
- Stage 5: sub-pixel search at interpolated positions.

SUMMARY OF THE INVENTION

A parallel motion estimation architecture is shown that enables efficient utilization of computational resources by making use of the inherent parallelism of the motion estimation algorithms.

BRIEF DESCRIPTION OF THE DRAWINGS

These and other aspects of this invention are illustrated in the drawings, in which:

FIG. 1 illustrates the organization of a typical digital signal processor to which this invention is applicable (prior art);

FIG. 2 illustrates details of a very long instruction word digital signal processor core suitable for use in FIG. 1 (prior art);

FIG. 3 illustrates the pipeline stages of the very long instruction word digital signal processor core illustrated in FIG. 2 (prior art);

FIG. 4 illustrates the instruction syntax of the very long instruction word digital signal processor core illustrated in FIG. 2 (prior art);

FIG. 5 illustrates an overview of the video encoding process of the prior art;

FIG. 6 illustrates an overview of the video decoding process of the prior art;

FIG. 7 illustrates an overview of the motion estimation engine of this invention; and

FIG. 8 illustrates one of the local buffers.

DETAILED DESCRIPTION OF PREFERRED EMBODIMENTS

FIG. 1 illustrates the organization of a typical digital signal processor system **100** to which this invention is applicable (prior art). Digital signal processor system **100** includes central processing unit core **110**. Central processing unit core **110** includes the data processing portion of digital signal processor system **100**. Central processing unit core **110** could be constructed as known in the art and would typically include a register file, an integer arithmetic logic unit, an integer multiplier and program flow control units. An example of an appropriate central processing unit core is described below in conjunction with FIGS. 2 to 4.

Digital signal processor system **100** includes a number of cache memories. FIG. 1 illustrates a pair of first level caches. Level one instruction cache (L1I) **121** stores instructions used by central processing unit core **110**. Central processing unit core **110** first attempts to access any instruction from level one instruction cache **121**. Level one data cache (L1D) **123** stores data used by central processing unit core **110**. Central processing unit core **110** first attempts to access any required data from level one data cache **123**. The two level one caches are backed by a level two unified cache (L2) **130**. In the event of a cache miss to level one instruction cache **121** or to level one data cache **123**, the requested instruction or data is sought from level two unified cache **130**. If the requested instruction or data is stored in level two unified cache **130**, then it is supplied to the requesting level one cache for supply to central processing unit core **110**. As is known in the art, the requested instruction or data may be simultaneously supplied to both the requesting cache and central processing unit core **110** to speed use.

Level two unified cache **130** is further coupled to higher level memory systems. Digital signal processor system **100** may be a part of a multiprocessor system. The other processors of the multiprocessor system are coupled to level two unified cache **130** via a transfer request bus **141** and a data transfer bus **143**. A direct memory access unit **150** provides the connection of digital signal processor system **100** to external memory **161** and external peripherals **169**.

FIG. 2 is a block diagram illustrating details of a digital signal processor integrated circuit **200** suitable but not

essential for use in this invention (prior art). The digital signal processor integrated circuit **200** includes central processing unit **1**, which is a 32-bit eight-way VLIW pipelined processor. Central processing unit **1** is coupled to level one instruction cache **121** included in digital signal processor integrated circuit **200**. Digital signal processor integrated circuit **200** also includes level one data cache **123**. Digital signal processor integrated circuit **200** also includes peripherals **4** to **9**. These peripherals preferably include an external memory interface (EMIF) **4** and a direct memory access (DMA) controller **5**. External memory interface (EMIF) **4** preferably supports access to supports synchronous and asynchronous SRAM and synchronous DRAM. Direct memory access (DMA) controller **5** preferably provides 2-channel auto-boot loading direct memory access. These peripherals include power-down logic **6**. Power-down logic **6** preferably can halt central processing unit activity, peripheral activity, and phase lock loop (PLL) clock synchronization activity to reduce power consumption. These peripherals also include host ports **7**, serial ports **8** and programmable timers **9**.

Central processing unit **1** has a 32-bit, byte addressable address space. Internal memory on the same integrated circuit is preferably organized in a data space including level one data cache **123** and a program space including level one instruction cache **121**. When off-chip memory is used, preferably these two spaces are unified into a single memory space via the external memory interface (EMIF) **4**.

Level one data cache **123** may be internally accessed by central processing unit **1** via two internal ports **3a** and **3b**. Each internal port **3a** and **3b** preferably has 32 bits of data and a 32-bit byte address reach. Level one instruction cache **121** may be internally accessed by central processing unit **1** via a single port **2a**. Port **2a** of level one instruction cache **121** preferably has an instruction-fetch width of 256 bits and a 30-bit word (four bytes) address, equivalent to a 32-bit byte address.

Central processing unit **1** includes program fetch unit **10**, instruction dispatch unit **11**, instruction decode unit **12** and two data paths **20** and **30**. First data path **20** includes four functional units designated L1 unit **22**, S1 unit **23**, M1 unit **24** and D1 unit **25** and 16 32-bit A registers forming register file **21**. Second data path **30** likewise includes four functional units designated L2 unit **32**, S2 unit **33**, M2 unit **34** and D2 unit **35** and 16 32-bit B registers forming register file **31**. The functional units of each data path access the corresponding register file for their operands. There are two cross paths **27** and **37** permitting access to one register in the opposite register file each pipeline stage. Central processing unit **1** includes control registers **13**, control logic **14**, and test logic **15**, emulation logic **16** and interrupt logic **17**.

Program fetch unit **10**, instruction dispatch unit **11** and instruction decode unit **12** recall instructions from level one instruction cache **121** and deliver up to eight 32-bit instructions to the functional units every instruction cycle. Processing occurs simultaneously in each of the two data paths **20** and **30**. As previously described each data path has four corresponding functional units (L, S, M and D) and a corresponding register file containing 16 32-bit registers. Each functional unit is controlled by a 32-bit instruction. The data paths are further described below. A control register file **13** provides the means to configure and control various processor operations.

FIG. 3 illustrates the pipeline stages **300** of digital signal processor core **110** (prior art). These pipeline stages are divided into three groups: fetch group **310**; decode group **320**; and execute group **330**. All instructions in the instruc-

tion set flow through the fetch, decode, and execute stages of the pipeline. Fetch group **310** has four phases for all instructions, and decode group **320** has two phases for all instructions. Execute group **330** requires a varying number of phases depending on the type of instruction.

The fetch phases of the fetch group **310** are: Program address generate phase **311** (PG); Program address send phase **312** (PS); Program access ready wait stage **313** (PW); and Program fetch packet receive stage **314** (PR). Digital signal processor core **110** uses a fetch packet (FP) of eight instructions. All eight of the instructions proceed through fetch group **310** together. During PG phase **311**, the program address is generated in program fetch unit **10**. During PS phase **312**, this program address is sent to memory. During PW phase **313**, the memory read occurs. Finally during PR phase **314**, the fetch packet is received at CPU **1**.

The decode phases of decode group **320** are: Instruction dispatch (DP) **321**; and Instruction decode (DC) **322**. During the DP phase **321**, the fetch packets are split into execute packets. Execute packets consist of one or more instructions which are coded to execute in parallel. During DP phase **322**, the instructions in an execute packet are assigned to the appropriate functional units. Also during DC phase **322**, the source registers, destination registers and associated paths are decoded for the execution of the instructions in the respective functional units.

The execute phases of the execute group **330** are: Execute 1 (E1) **331**; Execute 2 (E2) **332**; Execute 3 (E3) **333**; Execute 4 (E4) **334**; and Execute 5 (E5) **335**. Different types of instructions require different numbers of these phases to complete. These phases of the pipeline play an important role in understanding the device state at CPU cycle boundaries.

During E1 phase **331**, the conditions for the instructions are evaluated and operands are read for all instruction types. For load and store instructions, address generation is performed and address modifications are written to a register file. For branch instructions, branch fetch packet in PG phase **311** is affected. For all single-cycle instructions, the results are written to a register file. All single-cycle instructions complete during the E1 phase **331**.

During the E2 phase **332**, for load instructions, the address is sent to memory. For store instructions, the address and data are sent to memory. Single-cycle instructions that saturate results set the SAT bit in the control status register (CSR) if saturation occurs. For single cycle 16 by 16 multiply instructions, the results are written to a register file. For M unit non-multiply instructions, the results are written to a register file. All ordinary multiply unit instructions complete during E2 phase **332**.

During E3 phase **333**, data memory accesses are performed. Any multiply instruction that saturates results sets the SAT bit in the control status register (CSR) if saturation occurs. Store instructions complete during the E3 phase **333**.

During E4 phase **334**, for load instructions, data is brought to the CPU boundary. For multiply extension instructions, the results are written to a register file. Multiply extension instructions complete during the E4 phase **334**.

During E5 phase **335**, load instructions write data into a register. Load instructions complete during the E5 phase **335**.

FIG. 4 illustrates an example of the instruction coding of instructions used by digital signal processor core **110** (prior art). Each instruction consists of 32 bits and controls the operation of one of the eight functional units. The bit fields are defined as follows. The creg field (bits **29** to **31**) is the conditional register field. These bits identify whether the

instruction is conditional and identify the predicate register. The z bit (bit **28**) indicates whether the predication is based upon zero or not zero in the predicate register. If z=1, the test is for equality with zero. If z=0, the test is for nonzero. The case of creg=0 and z=0 is treated as always true to allow unconditional instruction execution. The creg field is encoded in the instruction opcode as shown in Table 1.

TABLE

Conditional Register	creg			z
	31	30	29	
Unconditional	0	0	0	0
Reserved	0	0	0	1
B0	0	0	1	z
B1	0	1	0	z
B2	0	1	1	z
A1	1	0	0	z
A2	1	0	1	z
A0	1	1	0	z
Reserved	1	1	1	x

Note that “z” in the z bit column refers to the zero/not zero comparison selection noted above and “x” is a don’t care state. This coding can only specify a subset of the 32 registers in each register file as predicate registers. This selection was made to preserve bits in the instruction coding.

The dst field (bits **23** to **27**) specifies one of the 32 registers in the corresponding register file as the destination of the instruction results.

The scr2 field (bits **18** to **22**) specifies one of the 32 registers in the corresponding register file as the second source operand.

The scr1/cst field (bits **13** to **17**) has several meanings depending on the instruction opcode field (bits **2** to **12**). The first meaning specifies one of the 32 registers of the corresponding register file as the first operand. The second meaning is a 5-bit immediate constant. Depending on the instruction type, this is treated as an unsigned integer and zero extended to 32 bits or is treated as a signed integer and sign extended to 32 bits. Lastly, this field can specify one of the 32 registers in the opposite register file if the instruction invokes one of the register file cross paths **27** or **37**.

The opcode field (bits **2** to **12**) specifies the type of instruction and designates appropriate instruction options. A detailed explanation of this field is beyond the scope of this invention except for the instruction options detailed below.

The s bit (bit **1**) designates the data path **20** or **30**. If s=0, then data path **20** is selected. This limits the functional unit to L1 unit **22**, S1 unit **23**, M1 unit **24** and D1 unit **25** and the corresponding register file A **21**. Similarly, s=1 selects data path **30** limiting the functional unit to L2 unit **32**, S2 unit **33**, M2 unit **34** and D2 unit **35** and the corresponding register file B **31**.

The p bit (bit **0**) marks the execute packets. The p-bit determines whether the instruction executes in parallel with the following instruction. The p-bits are scanned from lower to higher address. If p=1 for the current instruction, then the next instruction executes in parallel with the current instruction. If p=0 for the current instruction, then the next instruction executes in the cycle after the current instruction. All instructions executing in parallel constitute an execute packet. An execute packet can contain up to eight instructions. Each instruction in an execute packet must use a different functional unit.

FIG. 5 illustrates the encoding process **500** of video encoding according to the prior art. Many video encoding

standards use similar processes such as represented in FIG. 5. Encoding process **500** begins with the nth (current) frame F_n **501**. Frequency transform block **502** transforms a macroblock of the pixel data into the spatial frequency domain. This typically involves a discrete cosine transform (DCT). This frequency domain data is quantized in quantization block **503**. This quantization typically takes into account the range of data values for the current macroblock. Thus differing macroblocks may have differing quantizations. In accordance with the H.264 standard, in the base profile the macroblock data may be arbitrarily reordered via reorder block **504**. As will be explained below, this reordering is reversed upon decoding. Other video encoding standards and the H.264 main profile transmit data for the macroblocks in strict raster scan order. The quantized data is encoded by entropy encoding block **505**. Entropy encoding employs fewer bits to encode more frequently used symbols and more bits to encode less frequently used symbols. This process reduces the amount of encoded data that must be transmitted and/or stored. The resulting entropy encoded data is the encoded data stream. This invention concerns content adaptive binary arithmetic coding (CABAC) which will be further described below.

Video encoding standards typically permit two types of predictions. In inter-frame prediction, data is compared with data from the corresponding location of another frame. In intra-frame prediction, data is compared with data from another location in the same frame.

For inter prediction, data from n-1th (previous) frame F_{n-1} **510** and data from the nth frame F_n **501** supply motion estimation block **511**. Motion estimation block **511** determines the positions and motion vectors of moving objects within the picture. This motion data is supplied to motion compensation block **512** along with data from n-1th frame F_{n-1} **510**. The resulting motion compensated frame data is selected by switch **513** for application to subtraction unit **506**. Subtraction unit **506** subtracts the inter prediction data from switch **513** from the input frame data from nth frame F_n **501**. Thus frequency transform block **502**, quantization block **503**, reorder block **504** and entropy encoding block **505** encode the differential data rather than the original frame data. Assuming there is relatively little change from frame to frame, this differential data has a smaller magnitude than the raw frame data. Thus this can be expressed in fewer bits contributing to data compression. This is true even if motion estimation block **511** and motion compensation block **512** find no moving objects to code. If then th frame F_n and the n-1th frame F_{n-1} are identical, the subtraction unit **506** will produce a string of zeros for data. This data string can be encoded using few bits.

The second type of prediction is intra prediction. Intra prediction predicts a macroblock of the current frame from another macroblock of the current frame. Inverse quantization block **520** receives the quantized data from quantization block **503** and substantially recovers the original frequency domain data. Inverse frequency transform block **521** transforms the frequency domain data from inverse quantization block **520** back to the spatial domain. This spatial domain data supplies one input of addition unit **522**, whose function will be further described. Encoding process **500** includes choose intra predication unit **514** to determine whether to implement intra prediction. Choose intra prediction unit **514** receives data from nth frame F_n **501** and the output of addition unit **522**. Choose intra prediction unit **514** signals intra prediction intra predication unit **515**, which also receives the output of addition unit **522**. Switch **513** selects the intra prediction output for application to the subtraction

input of subtraction units **506** and an addition input of addition unit **522**. Intra prediction is based upon the recovered data from inverse quantization block **520** and inverse frequency transform block **521** in order to better match the processing at decoding. If the encoding used the original frame, there might be drift between these processes resulting in growing errors.

Video encoders typically periodically transmit unpredicted frames. In such an event the predicted frame is all 0's. Subtraction unit **506** thus produces data corresponding to the n-th frame F_n **501** data. Periodic unpredicted or I frames limits any drift between the transmitter coding and the receive decoding. In a video movie a scene change may produce such a large change between adjacent frames that differential coding provides little advantage. Video coding standards typically signal whether a frame is a predicted frame and the type of prediction in the transmitted data stream.

Encoding process **500** includes reconstruction of the frame based upon this recovered data. The output of addition unit **522** supplies deblock filter **523**. Deblock filter **523** smoothes artifacts created by the block and macroblock nature of the encoding process. The result is reconstructed frame F'_n **524**. As shown schematically in FIG. 5, this reconstructed frame F'_n **524** becomes the next reference frame F_{n-1} **510**.

FIG. 6 illustrates the corresponding decoding process **600**. Entropy decode unit **601** receives the encoded data stream. Entropy decode unit **601** recovers the symbols from the entropy encoding of entropy encoding unit **505**. This invention is applicable to CABAC decoding. Reorder unit **602** assembles the macroblocks in raster scan order reversing the reordering of reorder unit **504**. Inverse quantization block **603** receives the quantized data from reorder unit **602** and substantially recovers the original frequency domain data. Inverse frequency transform block **604** transforms the frequency domain data from inverse quantization block **603** back to the spatial domain. This spatial domain data supplies one input of addition unit **605**. The other input of addition input **605** comes from switch **609**. In inter prediction mode switch **609** selects the output of motion compensation unit **607**. Motion compensation unit **607** receives the reference frame F'_{n-1} **606** and applies the motion compensation computed by motion compensation unit **512** and transmitted in the encoded data stream.

Switch **609** may also select an intra prediction mode. The intra prediction is signaled in the encoded data stream. If this is selected, intra prediction unit **608** forms the predicted data from the output of adder **605** and then applies the intra prediction computed by intra prediction block **515** of the encoding process **500**. Addition unit **605** recovers the predicted frame. As previously discussed in conjunction with encoding, it is possible to transmit an unpredicted or I frame. If the data stream signals that a received frame is an I frame, then the predicted frame supplied to addition unit **605** is all 0's.

The output of addition unit **605** supplies the input of deblock filter **610**. Deblock filter **610** smoothes artifacts created by the block and macroblock nature of the encoding process. The result is reconstructed frame F''_n **611**. As shown schematically in FIG. 6, this reconstructed frame F''_n **611** becomes the next reference frame F_{n-1} **606**.

The deblocking filtering of deblock filter **523** and deblock filter **610** must be the same. This enables the decoding process to accurately reflect the input frame F_n **501** without error drift. The H.264 standard has a specific, very detailed decision matrix and corresponding filter operations for this

process. The standard deblock filtering is applied to every macroblock in raster scan order. This deblock filtering smooths artifacts created by the block and macroblock nature of the encoding. The filtered macroblock is used as the reference frame in predicted frames in both encoding and decoding. The encoding and decoding apply the identical processing the reconstructed frame to reduce the residual error after prediction.

The architecture shown in this invention enables efficient utilization of computational hardware by making use of the inherent parallelism in the motion estimation algorithm. Two types of parallelism are exploited: (1) For a given current macro-block, search operations in different reference frames are independent of each other and can be performed in parallel (or pipelined fashion). (2) Search operations for two different macro-blocks can be performed in parallel.

Motion estimation algorithms many times involve a global search followed by a local search phase. Global search narrows down the search to one (or more) motion vector candidates among several initial candidates. The local search that comes afterwards involves searching in a region just around the candidate(s) identified during the global search phase. Further, local search typically involves integer motion vector search and fractional (sub-pixel) motion vector search.

The architecture shown in FIG. 7 provides parallel hardware comprising of:

Interface controller **708** connected to processing unit **711**, skip engine **706**, skip input buffer **707**, I/O management unit **710**, predictor engine **701**, vector SAD engine **702** and subpel engine **703**.

Programming interface **709** is connected to skip engine **706** and skip buffer **705**.

I/O management unit **710** is connected to interface controller **708**, predictor engine **701**, vector SAD engine **702** and subpel engine **703**.

Processing unit **711** is connected to interface controller **708**, predictor engine **701**, vector SAD engine **702**, subpel engine **703** and interpolation reference buffer **704**.

Predictor engine **701** is connected to interface controller **708**, vector SAD engine **702** and subpel (sub pixel) engine **703**.

Subpel engine **703** is connected vector SAD engine **702**, predictor engine **701**, interpolation reference buffer **704**, processing unit **711** and interface controller **708**.

Vector SAD engine is connected to predictor engine **701**, subpel engine **703**, processing unit **711** and interface controller **708**.

Skip engine **706** is connected to interface controller **708**, skip input buffer **707**, skip MB buffer **705** and programming interface **709**.

Buffers **704**, **705** and **707** allow the implementation of global search, local integer search and local sub-pixel search operations in parallel. With this architecture, searches in multiple reference frames can also be performed in pipelined fashion, as shown in Table 2. This table shows L0 and L1 (forward and backward) direction search in B frame processing. After completing predictor search (global search) in L0 direction, local integer search (also called vector search) in L0 direction and global search in L1 direction happen in parallel. Similarly, local integer search in L1 direction happens in parallel with local sub pixel search in L direction.

A separate sub pixel interpolation and cost calculation engine **703** is used to calculate skip and direct mode costs as defined in H.264 standard.

9

TABLE 2

Predictor Search	Vector Search	Sub Pixel Search	Skip SAD
MBnL0 MBnL1	MBnL0 MBnL1	MBnL0 MBnL1	MBnL0 MBnL1
MBn + 1 L0 MBn + 1 L1	MBn + 1 L0 MBn + 1 L1	MBn + 1 L0 MBn + 1 L1	MBn + 1 L0 MBn + 1 L1

TABLE 3

Predictor Search	Vector Search	Sub Pixel Search	Skip Sad
MBn + 1 L0 MBn + 1 L1	MBnL0 MBnL1	MBnL0 MBnL1	MBnL0 MBnL1
MBn + 2 L0 MBn + 2 L1	MBn + 1 L0 MBn + 1 L1	MBn + 1 L0 MBn + 1 L1	MBn + 1 L0 MBn + 1 L1

For optimizing memory bandwidth, the architecture uses local buffers **704**, **705** and **707** to store the reference data required for portions of the search. The local buffers also free up memory ports and thus enable more parallel operations.

During local search, as shown in FIG. **8**, the same reference pixels are repeatedly accessed from the memory **801** by the search engine **804**. We take advantage of this fact and store all the data that may be required to perform the local search in custom designed local buffers **803** inside the search engine. Since the local search pattern is be known ahead of time, a bounded region of reference pixels required for local search around the starting point of local search (referred as LS_start_mv) can be identified.

After performing global search, before starting local search this bounded region is fetched from the memory and stored inside a buffer local to the motion estimation engine. Entire local search is then carried out by accessing reference data from the local buffer instead of fetching from the main memory. This significantly reduces the amount of data required to be fetched from the memory, thereby saving power and freeing up memory bandwidth.

More than one instance of local buffers may be included in the design to store more than one reference region required for local search. These reference regions may belong to either the same or to different reference frames.

In addition to storing the integer pixel reference data, the local buffer can also be used to store sub-pixel reference data. For example, H.264 standard specifies a 6-tap filter for half pixel calculation and 2-tap filter for quarter pixel calculation. For quarter pixel calculation, half position pixels are required. During the search operation, typically quarter pixel search is performed after half pixel search. The half pixels calculated during half pixel search can be stored in the local buffer so that during quarter pixel search, half pixel positions need not be recalculated.

The invention claimed is:

1. A device comprising:
 - a vector engine;
 - a sub-pixel engine; and
 - a predictor engine configured to:
 - during a first processing cycle, perform a first global search based on a first macroblock of a first frame to

10

identify a first set of candidate macroblocks of a first reference frame that correspond to the first macroblock; and

during a second processing cycle, perform a second global search based on the first macroblock to identify a second set of candidate macroblocks of a second reference frame that correspond to the first macroblock, wherein the vector engine is configured to:

during the second processing cycle, perform a first local integer search in a first direction based on the first set of candidate macroblocks in parallel with the second global search; and

during a third processing cycle, perform a second local integer search in a second direction based on the second set of candidate macroblocks, wherein the sub-pixel engine is configured to:

during the third processing cycle, perform a first local fractional search in the first direction based on the first set of candidate macroblocks in parallel with the second local integer search; and

during a fourth processing cycle, perform a second local fractional search in the second direction based on the second set of candidate macroblocks, and

wherein the device is configured to determine a motion vector for the first macroblock relative to the first set of candidate macroblocks or the second set of candidate macroblocks based on the first and second global searches, further based on the first and second local integer searches, and further based on the first and second local fraction searches,

Specifically, Chou, Schaub, or Murali do not teach the amended claims where the specific limitations of parallel four processing cycles, reciting inter alia; wherein the device is configured to determine a motion vector for the first macroblock relative to the first set of candidate macroblocks or the second set of candidate macroblocks based on the first and second global searches, further based on the first and second local integer searches, and further based on the first and second local fraction searches.

2. The device of claim **1**, further comprising a skip engine configured to:

during the second processing cycle, determine a first sum of absolute differences (SAD) value between the first macroblock and the first set of candidate macroblocks; and

during the third processing cycle, perform a second determination of a SAD value between the first macroblock and the second set of candidate macroblocks.

3. The device of claim **1**, wherein the first local fractional search comprises a half-pixel search.

4. The device of claim **1**, wherein the first local fractional search comprises a quarter-pixel search.

5. The device of claim **1**, wherein the first, second, third, and fourth processing cycles are consecutive cycles.

6. The device of claim **1**,

wherein to perform the first local integer search, the vector engine is configured to search a region of the first reference frame around the first set of candidate macroblocks, and

wherein to perform the second local integer search, the vector engine is configured to search a region of the second reference frame around the second set of candidate macroblocks.

7. The device of claim **1**, wherein to perform the first global search, the predictor engine is configured to identify

11

a third set of candidate macroblocks of a third reference frame that correspond to the first macroblock.

8. The device of claim 1, wherein the first reference frame is prior to the first frame.

9. The device of claim 1, wherein the predictor engine is further configured to perform a third global search based on a second macroblock of the first frame using an approximation of the motion vector for the first macroblock.

10. The device of claim 1, wherein to determine the motion vector, the device is configured to determine a best match based on a vector sum of absolute differences (SAD).

11. The device of claim 10, wherein the device is configured to determine the motion vector based on the SAD between the first macroblock and a first candidate macroblock or the SAD between the first macroblock and a second candidate macroblock.

12. A method comprising:

during a first processing cycle, performing a first global search based on a first macroblock of a first frame to identify a first set of candidate macroblocks of a first reference frame that correspond to the first macroblock; and

during a second processing cycle, performing a second global search based on the first macroblock to identify a second set of candidate macroblocks of a second reference frame that correspond to the first macroblock; during the second processing cycle, performing a first local integer search in a first direction based on the first set of candidate macroblocks in parallel with the second global search; and

during a third processing cycle, performing a second local integer search in a second direction based on the second set of candidate macroblocks,

during the third processing cycle, performing a first local fractional search in the first direction based on the first set of candidate macroblocks in parallel with the second local integer search; and

during a fourth processing cycle, performing a second local fractional search in the second direction based on the second set of candidate macroblocks, and

determining a motion vector for the first macroblock relative to the first set of candidate macroblocks or the second set of candidate macroblocks based on the first and second global searches, further based on the first and second local integer searches, and further based on the first and second local fraction searches.

13. The method of claim 12, further comprising:

during the second processing cycle, determine a first sum of absolute differences (SAD) value between the first macroblock and the first set of candidate macroblocks; and

during the third processing cycle, perform a second determination of a SAD value between the first macroblock and the second set of candidate macroblocks.

14. The method of claim 12, wherein the first local fractional search comprises a half-pixel search or a quarter-pixel search.

15. The method of claim 12, wherein the first, second, third, and fourth processing cycles are consecutive cycles.

16. The method of claim 12,

wherein performing the first local integer search comprises searching a region of the first reference frame around the first set of candidate macroblocks, and

wherein performing the second local integer search comprises searching a region of the second reference frame around the second set of candidate macroblocks.

12

17. The method of claim 12, wherein the method further comprises performing a third global search based on a second macroblock of the first frame using an approximation of the motion vector for the first macroblock.

18. The method of claim 12,

wherein determining the motion vector comprises determining a best match based on a vector sum of absolute differences (SAD), and

wherein determining the motion vector is based on the SAD between the first macroblock and a first candidate macroblock or the SAD between the first macroblock and a second candidate macroblock.

19. A device comprising:

a vector engine;

a sub-pixel engine; and

a predictor engine configured to:

during a first processing cycle, perform a first global search based on a first macroblock of a first frame to identify a first set of candidate macroblocks of a first reference frame that correspond to the first macroblock, wherein the first reference frame is prior to the first frame;

during a second processing cycle, perform a second global search based on the first macroblock to identify a second set of candidate macroblocks of a second reference frame that correspond to the first macroblock; and

perform a third global search based on a second macroblock of the first frame using an approximation of a motion vector for the first macroblock, wherein the vector engine is configured to:

during the second processing cycle, perform a first local integer search in a first direction based on the first set of candidate macroblocks in parallel with the second global search; and

during a third processing cycle, perform a second local integer search in a second direction based on the second set of candidate macroblocks,

wherein the sub-pixel engine is configured to:

during the third processing cycle, perform a first local fractional search in the first direction based on the first set of candidate macroblocks in parallel with the second local integer search; and

during a fourth processing cycle, perform a second local fractional search in the second direction based on the second set of candidate macroblocks, and

wherein the first, second, third, and fourth processing cycles are consecutive cycles, and

wherein the device is configured to determine the motion vector for the first macroblock relative to the first set of candidate macroblocks or the second set of candidate macroblocks based on the first and second global searches, further based on the first and second local integer searches, and further based on the first and second local fraction searches.

20. The device of claim 19, further comprising a skip engine configured to:

during the second processing cycle, determine a first sum of absolute differences (SAD) value between the first macroblock and the first set of candidate macroblocks; and

during the third processing cycle, perform a second determination of a SAD value between the first macroblock and the second set of candidate macroblocks.

* * * * *