

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent
Kind Code
Date of Patent
Inventor(s)

12394082
B2
August 19, 2025
Mckay; Jeremy et al.

Methods and systems for resolving 3D positions of multiple objects present within a 3D space

Abstract

An illustrative multi-view 3D positioning system assigns a respective 3D position estimate, with respect to a 3D space, to each of a plurality of detected instance datasets corresponding to objects present within the 3D space. Based on the 3D position estimates, the multi-view 3D positioning system sorts the plurality of detected instance datasets into a plurality of groupings corresponding to the objects so that each detected instance dataset is grouped together with other detected instance datasets corresponding to a same object. The multi-view 3D positioning system then determines a respective 3D position resolution, with respect to the 3D space, for each of the plurality of objects. These 3D position resolutions are determined based on the plurality of groupings of detected instance datasets and represent, with greater accuracy than the 3D position estimates, 3D positions of the objects within the 3D space. Corresponding methods and systems are also disclosed.

Inventors: Mckay; Jeremy (Englewood, CO), Hou; Tim KuanCheng (Denver, CO), Aldeborgh; Hannah Nicole (Wheat Ridge, CO), Haldimann; Ryan (Seattle, WA), Young; Brita Bunnell Schneiders (Denver, CO), MacMillan; Gordon David (Boulder, CO), Clague; James (Longmont, CO), Barnes; Jared David (Denver, CO), Panettiere; John Marco (Schnecksville, PA), Escanilla; Nicholas Sean (Broomfield, CO), McGuinness; William John (Golden, CO), Ren; Xiaolin (Queen Creek, AZ)

Applicant: Verizon Patent and Licensing Inc. (Basking Ridge, NJ)

Family ID: 1000008767176

Assignee: Verizon Patent and Licensing Inc. (Basking Ridge, NJ)

Appl. No.: 17/889012

Filed: August 16, 2022

Prior Publication Data

Document Identifier

Publication Date

Publication Classification

Int. Cl.: G06T7/70 (20170101); G06T7/292 (20170101); G06T7/50 (20170101); G06V10/764 (20220101); G06V20/50 (20220101)

U.S. Cl.:

CPC G06T7/70 (20170101); G06T7/292 (20170101); G06T7/50 (20170101); G06V10/764 (20220101); G06V20/50 (20220101); G06T2207/10016 (20130101)

Field of Classification Search

CPC: G06T (7/70); G06T (7/292); G06T (7/50); G06T (2207/10016); G06V (10/764); G06V (20/50); G06V (10/806); G06V (20/52)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
11315262	12/2021	Cherevatsky	N/A	H04N 7/188
2008/0056536	12/2007	Hildreth	382/103	G06V 40/107
2015/0199817	12/2014	Ishimaru	382/103	G01S 5/16
2021/0056297	12/2020	Shimshoni	N/A	G06V 40/103
2021/0183098	12/2020	Huang	N/A	G06T 7/596
2022/0108470	12/2021	Brown	N/A	G06T 7/75
2024/0282001	12/2023	Gonzalez	N/A	G06T 7/73

Primary Examiner: Mahrouka; Wassim

Background/Summary

BACKGROUND INFORMATION

(1) As an object moves within a 3D space, it may be useful for various applications to keep track of that object and its position within the space. Accordingly, given a video sequence depicting an object in a 3D space, various techniques and technologies have been developed to help achieve this object tracking. For example, certain technologies use computer vision to recognize and identify a particular object within a video sequence, and to track the object from frame to frame of the video as the object moves in the space.

(2) While such technologies are useful under many circumstances, they may be found lacking when it comes to more complex situations. For instance, various challenges may be encountered when trying to simultaneously track a plurality of different objects moving within a scene (including, perhaps, a large number of similar-looking objects or objects of the same type) that would not need to be addressed by technologies configured to track a single moving object in a scene. Additionally, certain 3D spaces may be too large to be adequately captured from a single viewpoint or by a single image capture device (e.g., video camera, etc.), thereby creating tracking challenges and

complexities as objects move in and out of different fields of view of different image capture devices that may or may not overlap. Circumstances characterizing certain applications may also introduce design requirements (e.g., timing-related requirements, indoor/outdoor usage requirements, cost requirements, ease of use/setup requirements, etc.) that existing technologies are not prepared to adequately satisfy.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The accompanying drawings illustrate various embodiments and are a part of the specification. The illustrated embodiments are merely examples and do not limit the scope of the disclosure. Throughout the drawings, identical or similar reference numbers designate identical or similar elements.

(2) FIG. 1 shows an illustrative multi-view 3D positioning system configured to resolve 3D positions of multiple objects present within a 3D space.

(3) FIG. 2 shows an illustrative method for resolving 3D positions of multiple objects present within a 3D space.

(4) FIG. 3 shows an illustrative 3D space including multiple objects that may be tracked by multi-view 3D positioning systems described herein.

(5) FIG. 4 shows an illustrative implementation of a multi-view 3D positioning system operating in an illustrative configuration.

(6) FIG. 5 shows an illustrative image depicting a 3D space with multiple objects and several illustrative algorithms used by a multi-view 3D positioning system to process the image.

(7) FIG. 6 shows illustrative aspects of a variety of detected instance datasets used by a multi-view 3D positioning system for resolving 3D positions of multiple objects present within a 3D space.

(8) FIG. 7 shows illustrative aspects of how a sorting facility of a multi-view 3D positioning system may sort detected instance datasets in furtherance of resolving 3D positions of multiple objects present within a 3D space.

(9) FIG. 8 shows illustrative aspects of how a multi-view 3D positioning system may determine a 3D position resolution for an object based on a grouping of detected instance datasets.

(10) FIG. 9 shows an illustrative point-of-interest data store used by a multi-view 3D positioning system for resolving 3D positions of multiple objects present within a 3D space.

(11) FIG. 10 shows an illustrative mapping of the 3D position resolutions of multiple objects within a 3D space.

(12) FIG. 11 shows an illustrative computing device that may implement multi-view 3D positioning systems and/or other computing systems described herein.

DETAILED DESCRIPTION OF EXAMPLE EMBODIMENTS

(13) Methods and systems for resolving 3D positions of multiple objects present within a 3D space are described herein. Such systems may be referred to as multi-view 3D positioning systems, since the system may determine the 3D positions for objects (e.g., first by roughly estimating the positions and then by resolving the positions more accurately as will be described) using imagery captured from multiple viewpoints within the 3D space (e.g., by different image capture devices positioned at different locations to provide views of different parts of the 3D space). Additionally, as further described herein, once 3D positions for objects have been determined, multi-view 3D positioning systems may track the objects within the 3D space as the objects move about the 3D space as time passes.

(14) In this way, useful information (e.g., analytics information, logistics information, etc.) may be obtained and provided to the advantage of a variety of different types of applications and use cases. For instance, one illustrative application or use case that will be referred to and illustrated in the

description below to provide context to various principles is a warehouse application. In this application, the 3D space will be understood to be a warehouse and the objects may include things such as inventory stored in the warehouse (e.g., boxes, packages, pallets, and/or other such items), personnel working in the warehouse, equipment used by the personnel (e.g., forklifts, handtrucks, etc.), and so forth.

(15) This warehouse application is convenient for illustrating various principles associated with methods and systems described herein, including that they may operate in real time, that they may operate indoors and within a relatively large area (e.g., an area that requires many image capture devices to fully capture), that the overall number of objects being tracked may be large, that the scene may be relatively complex, and so forth. That being said, it will also be understood that principles described herein in the context of the warehouse example may also be readily applied to a variety of other applications and use cases. Some of these additional applications and use cases may not require or implement real-time operation; may be associated with outdoor spaces or other types of indoor spaces (e.g., indoor spaces that are smaller than a warehouse and/or are used for different purposes); may involve different numbers and types of objects being located and tracked; may include different types, numbers, and setups for image capture devices and computing devices processing captured imagery; and so forth. For instance, a few additional (and non-limiting) examples of applications and use cases that may benefit from methods and systems described herein may include: providing real-time corrections for fleets of drones, managing shipping ports where large numbers of shipping containers are being moved, analyzing construction sites where structures are built by large numbers of workers, and tracking objects of interest in connection with automated guided vehicles (AGVs) moving through 3D space (e.g., tracking the objects precisely without requiring use of detectable measurement devices such as laser range finders, etc.). In these or other examples, image capture devices associated with ground-based and/or air-based imaging assets (e.g., planes, satellites, ground-based imaging, etc.) may be combined to locate and track objects in the air or on the ground.

(16) As will be described in more detail below, multi-view 3D positioning systems described herein may operate using computer vision techniques that allow for 3D positions of multiple objects to be resolved and tracked without reliance on depth detection techniques that are less cost-effective (e.g., techniques involving time-of-flight (ToF) principles, etc.), less flexible or effective (e.g., techniques involving structure-from-motion (SfM) principles requiring movement of the viewpoints themselves, etc.), or that suffer from other similar drawbacks. Rather, multi-view 3D positioning systems described herein may be implemented using equipment that is readily available (e.g., off-the-shelf image capture devices and embedded computing systems) and cost effective enough to allow for deployment of large numbers of devices throughout a 3D space (e.g., to allow for many cameras and computing devices to be placed throughout a large warehouse).

(17) Once deployed and calibrated, multi-view 3D positioning systems described herein may operate by performing any of the functions described herein. For example, image capture devices may capture imagery of a 3D space and various objects included therein. On-site computing devices associated with the image capture devices (e.g., including computing resources integrated into the image capture devices themselves) may be used to identify objects within the imagery and to track the objects in two dimensions from frame to frame, assigning contiguous identifiers to the objects. Based on pre-calibrated extrinsic parameters of the image capture devices, on-premise computing devices may use principles of monocular depth estimation to roughly project objects captured by different image capture devices into a 3D space shared by other image capture devices. On-premise or off-premise computing devices (e.g., including multi-access servers such as multi-access edge compute (MEC) servers implemented on provider networks) may analyze these projections to sort them into groupings of projections that represent the same physical object as captured from different viewpoints. Based on these groupings, the various estimations of the 3D positions of the object may then be resolved and refined to be more precise by triangulating the

projections within a grouping (resulting in 3D position resolutions of the objects). A registry of objects and their 3D position resolutions may be maintained and updated as additional data is received, and the multi-view 3D positioning system may publish and/or otherwise use information representative of the objects and their resolved locations to accomplish whatever objectives a particular application or use case may have. As one example, position information for various objects may be stored temporally in the registry such that a filtering algorithm (e.g., a Kalman filter, etc.) may be used to estimate velocity and/or acceleration of the objects through 3D space. (18) Along with flexibility and cost efficiency advantages that may arise from the equipment on which multi-view 3D positioning systems are implemented (as described above), other technical advantages and benefits may also be provided by systems and methods described herein. For example, object locating and tracking may be performed in real time or may be time shifted; multi-view 3D positioning systems may be readily scalable to include any number of views being used to locate and track any number of objects within an indoor or outdoor 3D space of arbitrary size; a desired accuracy of resolved 3D positions of objects may be readily balanced with cost/efficiency restraints by adding views to (or removing views from) the multi-view 3D positioning system (i.e., since, as will be made apparent below, more overlapping views capturing an object allows for the position of the object to be resolved more accurately); and so forth.

(19) Various specific implementations will now be described in detail with reference to the figures. It will be understood that the specific implementations described below are provided as non-limiting examples and may be applied in various situations. Additionally, it will be understood that other examples not explicitly described herein may also be captured by the scope of the claims set forth below. Methods and systems described herein for resolving 3D positions of multiple objects present within a 3D space may provide any of the benefits mentioned above, as well as various additional and/or alternative benefits that will be described and/or made apparent below.

(20) FIG. 1 shows an illustrative multi-view 3D positioning system **100** (“system **100**”) configured to resolve 3D positions of multiple objects present within a 3D space. System **100** may be implemented by computer resources such as processors, memory facilities, image capture devices (e.g., digital cameras, etc.), storage facilities, communication interfaces, and so forth. For example, system **100** may include or be communicatively coupled to a plurality of image capture devices that are disposed within the 3D space and are configured to capture 2D imagery on which various detected instance datasets described in more detail below may be based. Moreover, system **100** may further include memory storing instructions and one or more processors communicatively coupled to the memory and configured to execute the instructions to perform functions described herein. For instance, the memory and one or more processors may be implemented by one or more on-premise computing devices (e.g., relatively small and portable embedded systems devices) operating at a site of the 3D space, by an on-premise or off-premise server or other computing system communicatively coupled to the on-premise computing devices (e.g., a larger computing system configured to serve several or all of the on-premise computing devices) and configured to perform the same or other functions described herein, and/or within any other computing system as may serve a particular implementation.

(21) In some examples, on-premise computing devices may be communicatively coupled with the image capture devices and distributed in a one-to-one manner (i.e., such that there is a dedicated on-premise computing device for each individual image capture device). These on-premise computing devices may be configured to generate detected instance datasets based on 2D imagery captured by the image capture devices, to assign respective 3D position estimates to at least some of the detected instance datasets, and to perform other functions described herein. A server communicatively coupled to some or all of the on-premise computing devices may be implemented as a multi-access edge compute (MEC) server operating on a provider network (e.g., a carrier network provided by a cellular data services provider, etc.), by a cloud server accessible by way of the Internet, or the like. The server may be configured to sort detected instance datasets into

groupings corresponding to different objects, determine respective 3D position resolutions for objects based on the groupings of detected instance datasets, and perform other functions described herein.

(22) A generalized representation of system **100** is shown in FIG. **1** to include a memory **102** and a processor **104** selectively and communicatively coupled to one another. Memory **102** and processor **104** may each include or be implemented by computer hardware that is configured to store and/or execute computer software. Various other components of computer hardware and/or software not explicitly shown in FIG. **1** (e.g., image capture devices, etc.) may also be included within system **100**. In some examples, memory **102** and processor **104** may be distributed between multiple devices and/or multiple locations as may serve a particular implementation.

(23) Memory **102** may store and/or otherwise maintain executable data used by processor **104** to perform any of the functionality described herein. For example, memory **102** may store instructions **106** that may be executed by processor **104**. Memory **102** may be implemented by one or more memory or storage devices, including any memory or storage devices described herein, that are configured to store data in a transitory or non-transitory manner. Instructions **106** may be executed by processor **104** to cause system **100** to perform any of the functionality described herein.

Instructions **106** may be implemented by any suitable application, software, script, code, and/or other executable data instance. Additionally, memory **102** may also maintain any other data accessed, managed, used, and/or transmitted by processor **104** in a particular implementation.

(24) Processor **104** may be implemented by one or more computer processing devices, including general-purpose processors (e.g., central processing units (CPUs), graphics processing units (GPUs), microprocessors, etc.), special-purpose processors (e.g., application-specific integrated circuits (ASICs), field-programmable gate arrays (FPGAs), etc.), or the like. Using processor **104** (e.g., when processor **104** is directed to perform operations represented by instructions **106** stored in memory **102**), system **100** may perform functions associated with resolving 3D positions of multiple objects present within a 3D space in accordance with methods and systems described herein and/or as may serve a particular implementation.

(25) As one example of functionality that processor **104** may perform, FIG. **2** shows an illustrative method **200** for resolving 3D positions of multiple objects present within a 3D space. While FIG. **2** shows illustrative operations according to one implementation, other implementations may omit, add to, reorder, and/or modify any of the operations shown in FIG. **2**. In some examples, multiple operations shown in FIG. **2** or described in relation to FIG. **2** may be performed concurrently (e.g., in parallel) with one another, rather than being performed sequentially as illustrated and/or described. One or more of the operations shown in FIG. **2** may be performed by a multi-view 3D positioning system such as system **100** and/or any implementation thereof.

(26) In certain examples, operations of method **200** may be performed in real time so as to provide, receive, process, and/or use data described herein immediately as the data is generated, updated, changed, exchanged, or otherwise becomes available (e.g., determining accurate position data and tracking objects immediately as the objects move within the 3D space). In such examples, certain operations described herein may involve real-time data, real-time representations, real-time conditions, and/or other real-time circumstances. As used herein, “real time” will be understood to relate to data processing and/or other actions that are performed immediately, as well as conditions and/or circumstances that are accounted for as they exist in the moment when the processing or other actions are performed. For example, a real-time operation may refer to an operation that is performed immediately and without undue delay, even if it is not possible for there to be absolutely zero delay. Similarly, real-time data, real-time representations, real-time conditions, and so forth, will be understood to refer to data, representations, and conditions that relate to a present moment in time or a moment in time when decisions are being made and operations are being performed (e.g., even if after a short delay), such that the data, representations, conditions, and so forth are temporally relevant to the decisions being made and/or the operations being performed.

(27) Each of operations **202-206** of method **200** will now be described in more detail as the operations may be performed by an implementation of system **100** (e.g., by processor **104** executing instructions **106** stored in memory **102**).

(28) At operation **202**, system **100** may assign a respective 3D position estimate, with respect to a 3D space, to each of a plurality of detected instance datasets corresponding to a plurality of objects present within the 3D space. As will be described in more detail below, detected instance datasets may be generated (e.g., at operation **202** or prior to the performance of operation **202**) based on 2D imagery that is captured by a plurality of image capture devices that provide different views of a 3D space. As a simple example, for instance, a 3D space in which three objects are each captured by four cameras will be considered. For a particular moment in time (e.g., a period of time associated with a single frame of video captured by the cameras), up to twelve (i.e., 3 multiplied by 4) detected instance datasets may be generated to represent each of the three objects as captured from each of the four viewpoints of the cameras.

(29) Continuing this simple example, at operation **202**, some or all of these detected instance datasets may be assigned a 3D position estimate. For example, after the objects are recognized, identified, labeled, etc. (processes that will be described in more detail below), a monocular depth estimation technique may be performed with respect to each of the captured 2D images such that a relatively rough (i.e., imprecise) projection into a global coordinate space referenced by all of the cameras in the system may be made for each object. Because intrinsic and extrinsic parameters of each image capture device may be known with relatively high precision (e.g., based on a calibration process performed prior to operation **202**), the position of the object with respect to a particular camera may be determined relatively precisely in two dimensions (based on the 2D imagery), while the third dimension (i.e., the depth dimension) may involve a less precise estimation at this point in the process. To more precisely resolve this third dimension, method **200** may continue on to operations **204** and **206**, where 3D position estimates from multiple views for each object are combined.

(30) At operation **204**, system **100** may sort the plurality of detected instance datasets based on the 3D position estimates assigned to the detected instance datasets at operation **202**. For example, system **100** may sort the detected instance datasets into a plurality of groupings that each correspond to different objects of the plurality of objects so that each detected instance dataset is grouped together with other detected instance datasets corresponding to a same object. Referring to the simple example introduced above, for instance, operation **204** may involve sorting the twelve detected instance datasets into three groupings of four datasets each (one grouping for each of the three objects), based on the respective 3D position estimates. Each grouping in this example may correspond to a particular one of the objects and may include detected instance datasets for that object from each of the four cameras. Because the dimensions associated with greater certainty and the depth dimension associated with less certainty vary across the four detected instance datasets (due, in part, to differences between views from which the detected instance datasets are generated and also possibly due to other factors such as calibration accuracy, lighting, etc.), combining corresponding data within a particular grouping may allow for the 3D position to be refined and resolved into what is referred to herein as a 3D position resolution. It will be understood that 3D position resolutions are generally more accurate and precise than 3D position estimates due to the nature of how each is determined.

(31) At operation **206**, system **100** may determine a respective 3D position resolution, with respect to the 3D space, for each of the plurality of objects. For example, the 3D position resolutions for the various objects may be determined based on the plurality of groupings of detected instance datasets sorted at operation **204** by triangulation techniques and other techniques described herein. Ultimately, the twelve 3D position estimates of the simple example introduced above may be refined into three 3D position resolutions that represent, with greater accuracy than the 3D position estimates, respective 3D positions of each of the three example objects within the 3D space. This

resolved 3D position data may then be continuously tracked over time, provided to other systems, and/or otherwise used in any manner as may serve a particular implementation.

(32) To further illustrate certain principles mentioned in the description of method **200**, FIG. **3** shows an illustrative 3D space **302** that is outlined by a dashed square for purposes of illustration, though it will be understood that a 3D space may be any shape or size and may or may not be contained within a well-defined outer boundary. 3D space **302** is shown to include a plurality of objects **304** that may be tracked by a multi-view 3D positioning system described herein such as system **100**. While objects **304** are illustrated as differently sized circles in FIG. **3**, it will be understood that these objects may represent any suitable objects that may be present in a given 3D space. For example, referring back to the warehouse example mentioned above, 3D space **302** could be implemented as a warehouse space in which warehouse personnel use equipment to manage inventory stored in the warehouse space, and the plurality of objects **304** could then include one or more object associated with the warehouse personnel (e.g., a worker), one or more objects associated with the equipment (e.g., a forklift), one or more objects associated with the inventory stored in the warehouse space (e.g., a box), and so forth.

(33) FIG. **3** further shows a plurality of views **306** that could be used by an example multi-view 3D positioning system (e.g., an implementation of system **100**). As shown, views **306** are represented by arrows in the corners of 3D space **302** that are oriented to be facing inward toward 3D space **302** and the objects **304** included therein. These views **306** could be obtained, for example, by positioning four image capture devices (e.g., cameras) in the corners of 3D space **302** and pointing the devices inward toward the center of the space.

(34) FIG. **4** shows an illustrative implementation of system **100** operating in an illustrative configuration **400** to help fulfill objectives of a particular application (e.g., a warehouse application, any of the other applications or use cases mentioned herein, or any other suitable application or use case as may serve a particular implementation). As shown, this implementation of system **100** includes a plurality of image capture devices **402** associated one-to-one with a plurality of on-premise computing devices **404**. The on-premise computing devices **404** are communicatively coupled with a multi-access edge compute (MEC) server **406** that happens to serve all of the on-premise computing devices **404** in this example. MEC server **406** is shown to include various facilities implementing calibration processes **408**, a message broker **410**, position resolution services **412**, and a point-of-interest data store **414**. In configuration **400**, system **100** is shown to communicate by way of a network **416** with a user device **418** that implements, in this example, an application **420** and an analytics service **422**.

(35) Image capture devices **402**, on-premise computing devices **404**, and MEC server **406** may be implemented in accordance with devices and computing systems described above (e.g., cameras, small embedded systems devices, larger server devices, etc.). Network **416** may be implemented by any suitable private or public networks (e.g., a provider-specific wired or wireless communications network such as a cellular carrier network operated by a mobile carrier entity, a local area network (LAN), a wide area network, the Internet, etc.) and may use any communication technologies, devices, media, protocols, or the like, as may serve a particular implementation. User device **418** may represent any type of device where data provided by system **100** interfaces with some type of user to be applied for some type of use. For example, user device **418** may be a MEC server (e.g., the same MEC server **406** or another MEC server) or a cloud server accessed by a user or system with some interest in the location data provided by system **100**. As another example, user device **418** may be implemented by a user device such as a personal computer or mobile device (e.g., smartphone, tablet, etc.) by way of which a user makes use of the location data provided by system **100**.

(36) Specific functions and sub-elements of these major components of configuration **400** will now be described with reference both to FIG. **4** and to FIGS. **5-10**. Specifically, as indicated by tags with dotted outlines in FIG. **4**, FIG. **5** will be referenced to describe 2D imagery captured by image

capture devices **402** and how the 2D image data is processed, FIG. **6** will be referenced to illustrate detected instance datasets output by on-premise computing devices **404**, FIGS. **7** and **8** will be referenced to illustrate how position resolution services **412** may determine 3D position resolutions for various objects depicted in the 2D imagery, FIG. **9** will be referenced to illustrate aspects of an implementation of point-of-interest data store **414**, and FIG. **10** will be referenced to illustrate the output of system **100** that will be used by user device **418**.

(37) FIG. **5** shows an illustrative 2D image **502** depicting an example of 3D space **302** (an indoor warehouse space in this example) with multiple examples of objects **304** including a worker object **304-1**, a forklift object **304-2**, a box object **304-3** resting on the ground, another box object **304-4** that is being moved by the worker, and a handtruck object **304-5** being used by the worker to move box object **304-4**. It will be understood that the objects **304** depicted in 2D image **502** are provided only by way of example. Larger or more complex warehouse examples may thus include many more objects **304** and a larger 3D space **302**, while other examples may involve outdoor spaces or other types of spaces unrelated to warehouses.

(38) After being captured (e.g., by one of image capture devices **402**), 2D image **502** may be processed in various ways (e.g., by one of on-premise computing devices **404**). To illustrate, FIG. **5** shows three computer vision algorithms **504**, **506**, and **508** that may be used in the processing and analysis of 2D image **502** to thereby generate an estimated depth map **510** and to provide a plurality of detected instance datasets **512** and a plurality of 3D position estimates **514** for the objects **304** depicted within 2D image **502**. Each of these algorithms will now be described in more detail.

(39) Using an object recognition algorithm **504**, system **100** may be configured to identify each object **304** within 2D image **502**. More particularly, object recognition algorithm **504** may be used to semantically segment each object **304** from other objects and background content depicted in image **502**, and to identify a class in which each object is to be categorized. For example, object recognition algorithm **504** may segment the worker object **304-1** and classify him as a human being, segment the forklift object **304-2** and classify it as a vehicle (or, more specifically, as a forklift), segment each of the box objects **304-3** and **304-4** and classify them as inventory items (or, more specifically, as boxes), and so forth. In some implementations, object recognition algorithm **504** may be implemented by a You-Only-Look-Once (YOLO) algorithm or by another such object recognition algorithm (e.g., an established or custom algorithm). While FIG. **5** shows only one 2D image **502** depicting 3D space **302** from only one particular viewpoint, it will be understood that system **100** may use the same algorithm to identify these same objects **304** within other 2D images that depict 3D space **302** from other viewpoints (e.g., by being captured by image capture devices positioned elsewhere within the 3D space **302**). As will be made more apparent below, capturing data for each object **304** from multiple views **306** (e.g., due to overlapping fields of view of the image capture devices **402**) may allow for accurate 3D positioning to be performed by correlating the different detected instances of each object from the different viewpoints with one another and analyzing them in a manner that allows the multi-view information to be accumulated.

(40) To assist with such correlating, a tracking algorithm **506** may be used by system **100** to assign respective 2D identifiers to each particular object **304** identified within 2D image **502** by way of object recognition algorithm **504**. Tracking algorithm **506** may operate based on the identifying of objects **304** performed by object recognition algorithm **504** and may be configured to track object depictions that move from frame to frame in a 2D frame sequence. For example, the tracking algorithm may be implemented by a Deepsort tracking algorithm or another suitable tracking algorithm configured to label (e.g., assign 2D identifiers to) different objects in a manner that facilitates correlating the labeled object depictions with other labeled object depictions in other 2D images that are processed in the same way (and that depict the same objects **304**) at different points in time. As shown in FIG. **5**, and as will be described and illustrated in more detail below with reference to FIG. **6**, various detected instance datasets **512** may be generated for each instance of an

object being identified (e.g., by way of object recognition algorithm **504**) within a 2D image and labeled (e.g., by way of tracking algorithm **506**). Accordingly, tracking algorithm **506** may associate, within each detected instance dataset **512**, the particular object represented by the detected instance datasets with the respective 2D identifiers assigned to that particular object. As will be explained in more detail below, certain of these detected instance datasets **512** may be used straightaway (i.e., without waiting for a respective 3D position estimate **514**) while, for others detected instance datasets **512**, a 3D position estimate **514** may be generated and associated with the detected instance datasets **512**.

(41) 3D position estimates **514** may be generated and assigned to detected instance datasets **512** based on: 1) extrinsic parameters of the set of image capture devices capturing the 2D imagery (e.g., including 2D image **502**) on which the plurality of detected instance datasets **512** is based; and 2) an estimated depth map **510** for objects **304** depicted in the 2D imagery (where estimated depth map **510** is generated by a monocular depth estimation algorithm **508** configured to estimate depth based on monoscopic imagery). More particularly, as shown, 3D position estimates **514** are shown to be generated by way of a monocular data estimation algorithm **508** that may generate, maintain, update, and/or otherwise make use of an estimated depth map **510** of 3D space **302** from the viewpoint of 2D image **502**. Many depth detection algorithms require reliance on things like multiple 2D images from different viewpoints (e.g., stereoscopic depth detection algorithms), projections of light into the scene (e.g., structured light and time-of-flight depth detection algorithms), continuous movement of the image capture device viewpoint (e.g., structure-from-motion depth detection algorithms), and so forth. In contrast to these algorithms, monocular data estimation algorithm **508** may determine estimated depth map **510** based only on computer vision principles. More particularly, algorithm **508** may determine estimated depth map **510** based on the 2D image **502** itself and based on pre-calibrated extrinsics of the image capture device. For example, by applying a deep learning model trained up for the types of objects **304** that are found in a particular 3D space **302**, monocular data estimation algorithm **508** may approximate the depth of each object based on the appearance of the objects in the 2D image and based on what the model indicates for the estimated size of these objects (e.g., the average height of a male worker or a handtruck, the average length and width of a particular model of forklift, etc.).

(42) While estimated depth map **510** may not provide objective depth estimates (e.g., in feet or meters from the image capture device, etc.), it may nevertheless provide good enough relative approximations of the depths of the various objects that 3D position estimates **514** may be generated and assigned to each detected instance datasets **512** for use in the ways described herein (e.g., ultimately to facilitate sorting the detected instance datasets **512** and determining a highly accurate 3D position resolution for each object **304**).

(43) Any suitable component of system **100** may be configured to perform the functions of computer visions algorithms **504**, **506**, and **508**. For example, the image capture devices **402** may include sufficient onboard processing to run deep learning modules capable of implementing object recognition algorithm **504**, individual on-premise computing devices **404** corresponding to the image capture devices **402** in a one-to-one manner may then be configured to perform the functions of tracking algorithm **506** and monocular data estimation algorithm **508**, and the various detected instance datasets **512** and corresponding 3D position estimates **514** may then be provided to MEC server **406**.

(44) The plurality of detected instance datasets **512** may include various detected instance datasets that each correspond to a particular object **304** and are each generated based on the 2D image **502** depicting that particular object. It will be understood that other detected instance datasets corresponding to the same objects **304** may be generated based on other 2D images that depict the same objects **304** from different viewpoints. For example, an overall set of detected instance datasets **512** that may be provided by the set of on-premise computing devices **404** to MEC server **406** may include at least: 1) a first detected instance dataset **512** corresponding to a particular

object **304** of the plurality of objects, the first detected instance dataset generated based on a first 2D image **502** depicting the particular object **304** and captured by a first image capture device **402** in 3D space **302**; and 2) a second detected instance dataset **512** corresponding to that same particular object **304**, the second detected instance dataset generated based on a second 2D image **502** depicting the particular object **304** (i.e., a different 2D image from a different viewpoint) and captured by a second image capture device **402** in 3D space **302**. Moreover, various other detected instance datasets corresponding to this object **304**, other objects **304**, various 2D images captured from various viewpoints, etc., may also be included in the set provided to MEC server **406** in this example.

(45) To illustrate, FIG. **6** shows illustrative aspects of a variety of detected instance datasets **512** used by system **100** in this example. As shown, each detected instance dataset **512** represented in FIG. **6** is labeled using a numbering scheme that indicates both the object **304** and the view **306** that the detected instance dataset **512** corresponds to. Specifically, as used herein and as shown in FIG. **6**, a detected instance dataset **512-X-Y** will be understood to indicate that a detected instance of an object **304-Y** is represented as captured from a view **306-X**. For example, detected instance dataset **512-1-1** is shown to represent a detected instance dataset for object **304-1** as captured from an image capture device having view **306-1**, detected instance dataset **512-1-2** is shown to represent a detected instance dataset for object **304-2** as captured from the same image capture device with view **306-1**, detected instance dataset **512-2-1** is shown to represent a detected instance dataset for object **304-1** as captured from an image capture device having view **306-2**, and so forth. The respective objects and views are also indicated in the corner of each detected instance dataset **512** shown in FIG. **6**, as well as a frame number indication ("Frame **1**" in these examples).

(46) While the frame numbers for all of the detected instance datasets **512** shown in FIG. **6** are the same to indicate that they all represent approximately the same moment in time, it will be understood that a short time later (e.g., one frame time later, which may be a few tens of milliseconds or the like, depending on the frame rate of the video), a new set of detected instance datasets **512** (for Frame **2**, in this example) may be generated and analyzed in the same way as described herein for the set shown in FIG. **6**. It will also be understood that image capture devices **402** need not necessarily be synchronized, such that Frame **1** for one view **306** may not perfectly temporally line up with Frame **1** for another view **306**. Specifically, for example, the plurality of detected instance datasets **512** may include: 1) a first detected instance dataset **512** generated based on a first 2D image **502** that is captured by a first image capture device **402** in 3D space **302**, and 2) a second detected instance dataset **512** generated based on a second 2D image **502** that is captured, asynchronously with the first 2D image **502**, by a second image capture device **402** in 3D space **302**. A message buffer (e.g., a circular message buffer, etc.) implemented in the architecture receiving detected instance datasets **512** (e.g., within MEC server **406** in this example) may be used to analyze the detected instance datasets **512** that come in from the various image capture devices and on-premise computing devices even though they are not perfectly synchronized from view to view.

(47) As shown in FIG. **6**, each detected instance dataset **512** may include various datapoints (indicated with a monospace font within the detected instance datasets in FIG. **6**) indicative of various characteristics of the detected instance of the object **304** that is represented by the detected instance dataset. Specifically, as shown, each detected instance dataset **512** may include: a respective 2D identifier associated with the particular object **304**, a respective 2D position of a depiction of the particular object with respect to a 2D image **502** on which the detected instance dataset is based (e.g., in U-V coordinates from the particular view **306**), an indication of a size of the depiction of the particular object **304** with respect to the 2D image **502**, and any other indicators (e.g., the class of the object **304**, etc.) as may serve a particular implementation.

(48) FIG. **6** further shows, with dashed lines, that respective 3D position estimates **514** may be associated with at least some of the detected instance datasets **512** ("when available"). When a new

object appears in the 3D space (or in a portion of the 3D space that is visible from a particular view **306**), a 3D position estimate **514** may be determined and assigned to detected instance datasets **512** for that object. For example, this situation may occur at start up, when a new object is introduced to 3D space **302**, when an object that was previously occluded from a particular view becomes visible (e.g., something blocking the view moves to reveal what was behind it), or under other similar circumstances. While it would be suitable for every detected instance dataset **512** to be assigned a 3D position estimate **514**, this may not be necessary or even beneficial from a practical point of view. This is because monocular data estimation algorithm **508** may require significantly more computation than other algorithms involved in generating detected instance datasets **512**, such that it may not be practically possible to provide 3D position estimates **514** for every detected instance dataset **512** in real time. For example, while a frame rate for the detected instance datasets may be on the order of, e.g., 30 Hz, monocular data estimation algorithm **508** may only be able to process frames on the order of, e.g., 2 Hz. Fortunately, once a 3D position estimate **514** (and eventually a 3D position resolution) has been generated for a particular object viewed from a particular viewpoint, this information may be logged (described in more detail below) so that future detected instance datasets **512** associated with later frames do not necessarily need to be assigned respective 3D position estimates **514** to be useful. As such, once a point-of-interest dataset for a particular object **304** is generated, the point-of-interest dataset may be updated at a high rate (e.g., with every new video frame) without waiting for monocular depth estimates to be computed at the slower rate.

(49) Returning to FIG. 4, messages provided from each on-premise computing device **404** are shown to be received by MEC server **406** at message broker **410**. For example, these messages may include the detected instance datasets **512**, the 3D position estimates **514**, and/or other suitable data that has been described. Message broker **410** may be configured to use any technology to receive and route the large number of messages that may be received from on-premise computing devices **404** distributed throughout a 3D space. For example, message broker **410** may use an MQTT message technology or other suitable technology. In some examples, the 3D space managed by MEC server **406** may be large enough that messages may need to be routed to several position resolution services **412** associated with different parts of the 3D space (e.g., a northwest position resolution service for a northwest portion of a warehouse, a northeast position resolution service for a northeast portion of the warehouse, etc.).

(50) MEC server **406** is also shown to include calibration processes **408**, which will be understood to help initially determine and continually maintain accurate intrinsic and/or extrinsic parameters for each image capture device **402**. These parameters may be obtained by any suitable calibration processes **408** and, as has been mentioned and will further be made apparent, may be important for allowing system **100** to accurately resolve the 3D positions of the various objects present within the 3D space.

(51) Position resolution services **412** may receive (e.g., by way of routing processes in message broker **410**) relevant detected instance datasets **512** and/or associated 3D position estimates **514**, and, based on this received data, may be configured to both 1) sort the detected instance datasets into groupings corresponding to the various objects present at the scene (operation **204**), and 2) to determine the 3D position resolutions of the objects (operation **206**).

(52) To further illustrate these functions performed by position resolution services **412**, FIG. 7 shows illustrative aspects of how a sorting facility **702** of system **100** (e.g., implemented as part of position resolution service **412**) may sort detected instance datasets **512** in furtherance of resolving 3D positions of multiple objects present within a 3D space. In FIG. 7, detected instance datasets **512** are illustrated as small circles labeled in accordance with the same numbering scheme described above with respect to FIG. 6. Each of these detected instance datasets will be understood to be associated with a 3D position estimate that may be determined either according to the monocular depth estimation procedures described above (e.g., when the object is newly identified and its position not yet estimated) or by accessing data from a previous frame logged in a data store

(e.g., point-of-interest data store **414**, as described in more detail below). The small circles in FIG. **7** are positioned in accordance with the 3D position estimates for the detected instance datasets **512** and sorting facility **702** is shown to be configured to sort these detected instance datasets into different groupings **704** corresponding to different objects (e.g., grouping **704-1** for object **304-1**, grouping **704-2** for object **304-2**, and grouping **704-3** for object **304-3** in this example). As such, it is shown that each detected instance dataset **512** is grouped together with other detected instance datasets that correspond to a same object (e.g., all the detected instance datasets **512** corresponding to object **304-1** are grouped together into grouping **704-1**).

(53) Once the detected instance datasets **512** are correctly associated with one another by being clustered into their respective groupings **704** by sorting facility **702**, position calculations performed by position resolution services **412** may resolve an accurate 3D position for each object based on the various detected instance datasets **512** and corresponding 3D position estimates. For example, the determining of a particular 3D position resolution for a particular object **304** (operation **202**) may be performed based on each detected instance dataset **512** in a particular grouping **704** corresponding to that particular object **304**. This may be performed using any suitable technique (e.g., a triangulation technique, a singular value decomposition technique, etc.) in which a 3D position resolution of an object is determined to be a point of closest approach for rays from two or more viewpoints to where the object is detected from those respective viewpoints.

(54) To illustrate, FIG. **8** shows illustrative aspects of how system **100** (e.g., position resolution services **412**) may determine a 3D position resolution **802** for an object based on a grouping of detected instance datasets. As shown, various 2D images **502** (e.g., 2D images **502-1** through **502-4** in this example) that are captured from different views **306** (e.g., views **306-1** through **306-4**, respectively, in this example) may each include its own depiction **804** of a same object (e.g., a particular one of objects **304**, represented by a small black circle in this illustration). As shown, for instance, a depiction **804-1** of the object is included in 2D image **502-1**, a depiction **804-2** of the object is included in 2D image **502-2**, and so forth for depictions **804-3** and **804-4**. Depictions **804** are represented in respective detected instance datasets **512** that have been put into a same grouping **704** for the object in the ways described above. Due to small errors or imperfections in image capture devices, intrinsic and/or extrinsic calibration parameters, information in the detected instance datasets **512**, and/or other aspects, FIG. **8** shows that respective 3D position estimates **806** for the object **304** represented by depictions **804** (e.g., 3D position estimate **806-1** for depiction **804-1**, 3D position estimate **806-2** for depiction **804-2**, etc.) may be near one another, but not perfectly aligned. Accordingly, system **100** may determine 3D position resolution **802** to be disposed at a point that comes closest to all the rays overlapping (e.g., rays from views **306** through depictions **804**, represented by dotted lines in FIG. **8**). For instance, as mentioned above, this point may be calculated using a singular value decomposition technique, an averaging technique, or any other way of combining 3D position estimates **806** as may serve a particular implementation to generate 3D position resolution **802**.

(55) Returning to FIG. **4**, objects for which 3D position resolutions have been computed (e.g., by way of position resolution services **412** as described above) may be represented as point-of-interest datasets in a point-of-interest data store **414**. For example, once system **100** has completed the operations of method **200** and ultimately determined a 3D position resolution for a particular object present within a 3D space, system **100** may proceed to store, within point-of-interest data store **414**, a point-of-interest dataset for that object based on the 3D position resolution that was determined. In this way, system **100** may accumulate and update (track), within point-of-interest data store **414**, respective point-of-interest datasets for each object that comes into the 3D space.

(56) Each point-of-interest dataset stored in point-of-interest data store **414** may include certain data to represent the object and its 3D location as the object is tracked through the 3D space over time. For example, a particular point-of-interest dataset for a particular object may include a 3D position resolution (e.g., such as 3D position resolution **802**) determined for the particular object, a

set of detected instance datasets **512** that have been determined to correspond to the particular object, and/or any other suitable data as may serve a particular implementation. For example, a filtering technique (e.g., using a particle filter, a Kalman filter, etc.) may be used to track the object through the 3D space and positional information may be further smoothed and/or refined so that information such as velocity and acceleration may also be computed and stored.

(57) To illustrate, FIG. **9** shows an example implementation of point-of-interest data store **414** that may be used by system **100**. As shown, point-of-interest data store **414** includes a plurality of point-of-interest datasets **902** (e.g., point-of-interest datasets **902-1** through **902-4**, plus, potentially, various other datasets represented by an ellipsis but not explicitly shown in FIG. **9**). Each point-of-interest dataset **902** is associated with a different one of objects **304**. For instance, as parenthetically indicated in FIG. **9**, point-of-interest dataset **902-1** corresponds to object **304-1**, point-of-interest dataset **902-2** corresponds to object **304-2**, and so forth. In each point-of-interest dataset **902**, a respective 3D position resolution **904** (e.g., 3D position resolution **904-1** for point-of-interest dataset **902-1**, 3D position resolution **904-2** for point-of-interest dataset **902-2**, etc.) and a respective set of detected instance datasets **512** are also stored. Accordingly, for instance, system **100** may represent object **304-1** within point-of-interest dataset **902-1**, and may continuously update the detected instance datasets **512** that are detected for that object (e.g., permanently or on a rolling basis), as well as the resolved 3D position of object **304-1** (3D position resolution **904-1**) as object **304-1** moves in time.

(58) Whenever a new video frame is captured and analyzed to identify the objects depicted therein, new detected instance datasets may be registered to existing point-of-interest datasets **902** or new point-of-interest datasets **902** may be generated and added to point-of-interest data store **414** as needed. Accordingly, for a given set of detected instance datasets **512** that correspond to a particular object **304** and are included in a particular point-of-interest dataset **902**, the set may comprise: 1) one or more detected instance datasets from among the plurality of detected instance datasets to which the respective 3D positions estimates have been assigned; 2) one or more detected instance datasets that are not from among the plurality of detected instance datasets to which the respective 3D positions estimates have been assigned; and/or a combination of both types of detected instance datasets. In other words, the set of detected instance datasets **512** for a given point-of-interest dataset may include detected instance datasets **512** for “new” or “unknown” objects and to which respective 3D position estimates **514** were assigned (as described above in relation to FIG. **6**), detected instance datasets **512** for “previously registered” or “known” objects and to which respective 3D position estimates **514** have not been assigned, and/or a combination of both.

(59) Returning to FIG. **4**, system **100** may provide output representing the various objects (points of interest) represented in point-of-interest data store **414** to user device **418** for use in furtherance of any of the applications or use cases described herein or as may serve a particular implementation. For instance, in the warehouse example that has been used throughout the description above, application **420** may allow a user to search for any particular inventory item, personnel, or equipment and immediately discover where that object is located within the warehouse. As another example, a mapping showing some or all of the tracked items may be presented to allow the user to see, in real time, how people, equipment, and inventory are moving through the warehouse.

(60) FIG. **10** shows an example of such a mapping of the 3D position resolutions of multiple objects within a 3D space. Specifically, as shown, a mapping **1000** uses shaded ‘+’ symbols to indicate where the 3D position resolutions **904** for each of the objects is located within 3D space **302** (from an overhead view of the 3D space in this example). As indicated by a key **1002**, white ‘x’ symbols are also indicated in FIG. **10** to illustrate respective true 3D positions **1004** (e.g., true position **1004-1** of object **304-1**, true position **1004-2** of object **304-2**, etc.) of the respective objects in 3D space **302**. It will be understood that this ground truth information illustrated by true

locations **1004** may not actually be known to system **100** or provided to user device **418** (since the 3D position resolutions **904** represent the most accurate estimations of the positions known to system **100**). However, as can be seen in FIG. **10**, the calculated 3D position resolutions **904** may be perfectly or very closely aligned with ground truth (i.e., true positions **1004**), thereby providing very helpful and reliable information to application **420**.

(61) Returning to FIG. **4**, also shown within user device **418** along with application **420** is an analytics service **422**. Along with presenting mappings such as mapping **1000** and performing other application functions described above, user device **418** may also use the output data provided by system **100** to perform analytics to help, for example, make warehouse operations more efficient. For example, analytics service **422** may determine that inventory in certain parts of the warehouse is moved less efficiently, which may lead to a discovery that there is a crack in the floor that forklifts are needing to steer around and that analytics service **422** may recommend repairing to increase productivity. As another example, analytics service **422** may determine that one particular worker is especially effective and hardworking, and may recommend that the worker should be recognized and awarded accordingly. Various other such analytics may be performed by analytics service **422** for the warehouse application and/or for other applications and use cases as may serve a particular implementation.

(62) In certain embodiments, one or more of the processes described herein may be implemented at least in part as instructions embodied in a non-transitory computer-readable medium and executable by one or more computing devices. In general, a processor (e.g., a microprocessor) receives instructions, from a non-transitory computer-readable medium (e.g., a memory, etc.), and executes those instructions, thereby performing one or more operations such as the operations described herein. Such instructions may be stored and/or transmitted using any of a variety of known computer-readable media.

(63) A computer-readable medium (also referred to as a processor-readable medium) includes any non-transitory medium that participates in providing data (e.g., instructions) that may be read by a computer (e.g., by a processor of a computer). Such a medium may take many forms, including, but not limited to, non-volatile media, and/or volatile media. Non-volatile media may include, for example, optical or magnetic disks and other persistent memory. Volatile media may include, for example, dynamic random-access memory (DRAM), which typically constitutes a main memory. Common forms of computer-readable media include, for example, a disk, hard disk, magnetic tape, any other magnetic medium, a compact disc read-only memory (CD-ROM), a digital video disc (DVD), any other optical medium, random access memory (RAM), programmable read-only memory (PROM), electrically erasable programmable read-only memory (EEPROM), FLASH-EEPROM, any other memory chip or cartridge, or any other tangible medium from which a computer can read.

(64) FIG. **11** shows an illustrative computing device **1100** that may implement multi-view 3D positioning systems and/or other computing systems described herein. For example, computing device **1100** may include or implement (or partially implement) a multi-view 3D positioning system such as system **100**, an image capture device such as image capture devices **402**, an on-premise computing device such as on-premise computing devices **404**, a multi-access server such as MEC server **406** or a cloud server, a user device such as user device **418**, certain elements of a network such as network **416**, and/or any other computing devices or systems described herein (or any elements or subsystems thereof).

(65) As shown in FIG. **11**, computing device **1100** may include a communication interface **1102**, a processor **1104**, a storage device **1106**, and an input/output (I/O) module **1108** communicatively connected via a communication infrastructure **1110**. While an illustrative computing device **1100** is shown in FIG. **11**, the components illustrated in FIG. **11** are not intended to be limiting. Additional or alternative components may be used in other embodiments. Components of computing device **1100** shown in FIG. **11** will now be described in additional detail.

- (66) Communication interface **1102** may be configured to communicate with one or more computing devices. Examples of communication interface **1102** include, without limitation, a wired network interface (such as a network interface card), a wireless network interface (such as a wireless network interface card), a modem, an audio/video connection, and any other suitable interface.
- (67) Processor **1104** generally represents any type or form of processing unit capable of processing data or interpreting, executing, and/or directing execution of one or more of the instructions, processes, and/or operations described herein. Processor **1104** may direct execution of operations in accordance with one or more applications **1112** or other computer-executable instructions such as may be stored in storage device **1106** or another computer-readable medium.
- (68) Storage device **1106** may include one or more data storage media, devices, or configurations and may employ any type, form, and combination of data storage media and/or device. For example, storage device **1106** may include, but is not limited to, a hard drive, network drive, flash drive, magnetic disc, optical disc, RAM, dynamic RAM, other non-volatile and/or volatile data storage units, or a combination or sub-combination thereof. Electronic data, including data described herein, may be temporarily and/or permanently stored in storage device **1106**. For example, data representative of one or more executable applications **1112** configured to direct processor **1104** to perform any of the operations described herein may be stored within storage device **1106**. In some examples, data may be arranged in one or more databases residing within storage device **1106**.
- (69) I/O module **1108** may include one or more I/O modules configured to receive user input and provide user output. One or more I/O modules may be used to receive input for a single virtual experience. I/O module **1108** may include any hardware, firmware, software, or combination thereof supportive of input and output capabilities. For example, I/O module **1108** may include hardware and/or software for capturing user input, including, but not limited to, a keyboard or keypad, a touchscreen component (e.g., touchscreen display), a receiver (e.g., an RF or infrared receiver), motion sensors, and/or one or more input buttons.
- (70) I/O module **1108** may include one or more devices for presenting output to a user, including, but not limited to, a graphics engine, a display (e.g., a display screen), one or more output drivers (e.g., display drivers), one or more audio speakers, and one or more audio drivers. In certain embodiments, I/O module **1108** is configured to provide graphical data to a display for presentation to a user. The graphical data may be representative of one or more graphical user interfaces and/or any other graphical content as may serve a particular implementation.
- (71) In some examples, any of the facilities described herein may be implemented by or within one or more components of computing device **1100**. For example, one or more applications **1112** residing within storage device **1106** may be configured to direct processor **1104** to perform one or more processes or functions associated with processor **104** of system **100**. Likewise, memory **102** of system **100** may be implemented by or within storage device **1106**.
- (72) To the extent the aforementioned implementations collect, store, or employ personal information of individuals, groups or other entities, it should be understood that such information shall be used in accordance with all applicable laws concerning protection of personal information. Additionally, the collection, storage, and use of such information can be subject to consent of the individual to such activity, for example, through well known “opt-in” or “opt-out” processes as can be appropriate for the situation and type of information. Storage and use of personal information can be in an appropriately secure manner reflective of the type of information, for example, through various access control, encryption and anonymization techniques for particularly sensitive information.
- (73) In the preceding description, various example embodiments have been described with reference to the accompanying drawings. It will, however, be evident that various modifications and changes may be made thereto, and additional embodiments may be implemented, without

departing from the broader scope of the invention as set forth in the claims that follow. The specification and drawings are accordingly to be regarded in an illustrative rather than restrictive sense.

Claims

1. A method comprising: assigning a respective 3D position estimate, with respect to a 3D space, to each of a plurality of detected instance datasets corresponding to a plurality of objects present within the 3D space; sorting, based on the 3D position estimates, the plurality of detected instance datasets into a plurality of groupings corresponding to the plurality of objects so that each detected instance dataset is grouped together with other detected instance datasets corresponding to a same object; determining a respective 3D position resolution, with respect to the 3D space, for each of the plurality of objects, the 3D position resolutions determined based on the plurality of groupings of detected instance datasets and representing, with greater accuracy than the 3D position estimates, 3D positions of the plurality of objects within the 3D space; and storing, within a data store and based on the respective 3D position resolutions determined for each of the plurality of objects, respective point-of-interest datasets for each of the plurality of objects, wherein a particular point-of-interest dataset for a particular object includes: a 3D position resolution determined for the particular object; and a set of detected instance datasets that correspond to the particular object.
2. The method of claim 1, wherein: the plurality of detected instance datasets includes: a first detected instance dataset corresponding to a particular object of the plurality of objects, the first detected instance dataset generated based on a first 2D image depicting the particular object and captured by a first image capture device in the 3D space, and a second detected instance dataset corresponding to the particular object, the second detected instance dataset generated based on a second 2D image depicting the particular object and captured by a second image capture device in the 3D space; the method further comprises: identifying, using an object recognition algorithm, the particular object within the first 2D image and within the second 2D image, assigning, based on the identifying and using a tracking algorithm configured to track object depictions that move from frame to frame in a 2D frame sequence, respective 2D identifiers to the particular object for each of the respective first and second 2D images, and associating, within both the first and second detected instance datasets, the particular object with the respective 2D identifiers assigned to the particular object; and the sorting of the plurality of detected instance datasets into the plurality of groupings includes sorting, based on the respective 2D identifiers associated with the particular object within the first and second detected instance datasets, the first and second detected instance datasets into a same grouping that corresponds to the particular object.
3. The method of claim 2, wherein: the object recognition algorithm is implemented by a You-Only-Look-Once (YOLO) algorithm; and the tracking algorithm is implemented by a Deepsort tracking algorithm.
4. The method of claim 1, wherein the assigning of the respective 3D position estimates is based on: extrinsic parameters of a set of image capture devices that capture 2D imagery on which the plurality of detected instance datasets is based; and an estimated depth map for objects depicted in the 2D imagery, the estimated depth map generated by a monocular depth estimation algorithm configured to estimate depth based on monoscopic imagery.
5. The method of claim 1, wherein each detected instance dataset included in the particular point-of-interest dataset includes: a respective 2D identifier associated with the particular object; a respective 2D position of a depiction of the particular object with respect to a 2D image on which the particular detected instance dataset is based; and an indication of a size of the depiction of the particular object with respect to the 2D image.
6. The method of claim 1, wherein the set of detected instance datasets that correspond to the particular object and are included in the particular point-of-interest dataset comprises a detected

instance dataset from among the plurality of detected instance datasets to which the respective 3D positions estimates have been assigned.

7. The method of claim 1, wherein the set of detected instance datasets that correspond to the particular object and are included in the particular point-of-interest dataset comprises a detected instance dataset that is not from among the plurality of detected instance datasets to which the respective 3D positions estimates have been assigned.

8. The method of claim 1, wherein the plurality of detected instance datasets includes: a first detected instance dataset generated based on a first 2D image that is captured by a first image capture device in the 3D space, and a second detected instance dataset generated based on a second 2D image that is captured, asynchronously with the first 2D image, by a second image capture device in the 3D space.

9. The method of claim 1, wherein the determining of a particular 3D position resolution for a particular object is performed based on each detected instance dataset in a particular grouping corresponding to the particular object and including at least two detected instance datasets corresponding to the particular object.

10. The method of claim 1, wherein: the 3D space includes a warehouse space in which warehouse personnel use equipment to manage inventory stored in the warehouse space; and the plurality of objects includes at least one object associated with the warehouse personnel, at least one object associated with the equipment, and at least one object associated with the inventory stored in the warehouse space.

11. A system comprising: a memory storing instructions; and one or more processors communicatively coupled to the memory and configured to execute the instructions to perform a process comprising: assigning a respective 3D position estimate, with respect to a 3D space, to each of a plurality of detected instance datasets corresponding to a plurality of objects present within the 3D space; sorting, based on the 3D position estimates, the plurality of detected instance datasets into a plurality of groupings corresponding to the plurality of objects so that each detected instance dataset is grouped together with other detected instance datasets corresponding to a same object; determining a respective 3D position resolution, with respect to the 3D space, for each of the plurality of objects, the 3D position resolutions determined based on the plurality of groupings of detected instance datasets and representing, with greater accuracy than the 3D position estimates, 3D positions of the plurality of objects within the 3D space; and storing, within a data store and based on the respective 3D position resolutions determined for each of the plurality of objects, respective point-of-interest datasets for each of the plurality of objects, wherein a particular point-of-interest dataset for a particular object includes: a 3D position resolution determined for the particular object; and a set of detected instance datasets that correspond to the particular object.

12. The system of claim 11, further comprising a plurality of image capture devices disposed within the 3D space and configured to capture 2D imagery on which the plurality of detected instance datasets are based; wherein the memory and the one or more processors are implemented within: a plurality of on-premise computing devices operating at a site of the 3D space and communicatively coupled with the plurality of image capture devices, the plurality of on-premise computing devices configured to generate the detected instance datasets based on the 2D imagery and to perform the assigning of the respective 3D positions estimates; and a multi-access edge compute (MEC) server communicatively coupled to the plurality of on-premise computing devices and configured to perform the sorting and the determining of the respective 3D position resolutions based on the plurality of detected instance datasets and the respective 3D position estimates.

13. The system of claim 11, wherein: the plurality of detected instance datasets includes: a first detected instance dataset corresponding to a particular object of the plurality of objects, the first detected instance dataset generated based on a first 2D image depicting the particular object and captured by a first image capture device in the 3D space, and a second detected instance dataset corresponding to the particular object, the second detected instance dataset generated based on a

second 2D image depicting the particular object and captured by a second image capture device in the 3D space; the process further comprises: identifying, using an object recognition algorithm, the particular object within the first 2D image and within the second 2D image, assigning, based on the identifying and using a tracking algorithm configured to track object depictions that move from frame to frame in a 2D frame sequence, respective 2D identifiers to the particular object for each of the respective first and second 2D images, and associating, within both the first and second detected instance datasets, the particular object with the respective 2D identifiers assigned to the particular object; and the sorting of the plurality of detected instance datasets into the plurality of groupings includes sorting, based on the respective 2D identifiers associated with the particular object within the first and second detected instance datasets, the first and second detected instance datasets into a same grouping that corresponds to the particular object.

14. The system of claim 11, wherein the assigning of the respective 3D position estimates is based on: extrinsic parameters of a set of image capture devices that capture 2D imagery on which the plurality of detected instance datasets is based; and an estimated depth map for objects depicted in the 2D imagery, the estimated depth map generated by a monocular depth estimation algorithm configured to estimate depth based on monoscopic imagery.

15. A non-transitory computer-readable medium storing instructions that, when executed, direct a processor of a computing device to perform a process comprising: assigning a respective 3D position estimate, with respect to a 3D space, to each of a plurality of detected instance datasets corresponding to a plurality of objects present within the 3D space; sorting, based on the 3D position estimates, the plurality of detected instance datasets into a plurality of groupings corresponding to the plurality of objects so that each detected instance dataset is grouped together with other detected instance datasets corresponding to a same object; determining a respective 3D position resolution, with respect to the 3D space, for each of the plurality of objects, the 3D position resolutions determined based on the plurality of groupings of detected instance datasets and representing, with greater accuracy than the 3D position estimates, 3D positions of the plurality of objects within the 3D space; and storing, within a data store and based on the respective 3D position resolutions determined for each of the plurality of objects, respective point-of-interest datasets for each of the plurality of objects, wherein a particular point-of-interest dataset for a particular object includes: a 3D position resolution determined for the particular object; and a set of detected instance datasets that correspond to the particular object.

16. The non-transitory computer-readable medium of claim 15, wherein: the plurality of detected instance datasets includes: a first detected instance dataset corresponding to a particular object of the plurality of objects, the first detected instance dataset generated based on a first 2D image depicting the particular object and captured by a first image capture device in the 3D space, and a second detected instance dataset corresponding to the particular object, the second detected instance dataset generated based on a second 2D image depicting the particular object and captured by a second image capture device in the 3D space; the process further comprises: identifying, using an object recognition algorithm, the particular object within the first 2D image and within the second 2D image, assigning, based on the identifying and using a tracking algorithm configured to track object depictions that move from frame to frame in a 2D frame sequence, respective 2D identifiers to the particular object for each of the respective first and second 2D images, and associating, within both the first and second detected instance datasets, the particular object with the respective 2D identifiers assigned to the particular object; and the sorting of the plurality of detected instance datasets into the plurality of groupings includes sorting, based on the respective 2D identifiers associated with the particular object within the first and second detected instance datasets, the first and second detected instance datasets into a same grouping that corresponds to the particular object.

17. The non-transitory computer-readable medium of claim 15, wherein the assigning of the respective 3D position estimates is based on: extrinsic parameters of a set of image capture devices

that capture 2D imagery on which the plurality of detected instance datasets is based; and an estimated depth map for objects depicted in the 2D imagery, the estimated depth map generated by a monocular depth estimation algorithm configured to estimate depth based on monoscopic imagery.
