



US012394157B2

(12) **United States Patent**
Drummond et al.

(10) **Patent No.:** **US 12,394,157 B2**

(45) **Date of Patent:** Aug. 19, 2025

(54) **PLANNER FOR AN
OBJECTIVE-EFFECTUATOR**

(71) Applicant: **Apple Inc.**, Cupertino, CA (US)

(72) Inventors: **Mark Drummond**, Palo Alto, CA (US); **Siva Chandra Mouli Sivapurapu**, Santa Clara, CA (US); **Bo Morgan**, Emerald Hills, CA (US)

(73) Assignee: **Apple Inc.**, Cupertino, CA (US)

(*) Notice: Subject to any disclaimer, the term of this patent is extended or adjusted under 35 U.S.C. 154(b) by 81 days.

(21) Appl. No.: 17/685,471

(22) Filed: **Mar. 3, 2022**

(65) **Prior Publication Data**

US 2022/0262081 A1 Aug. 18, 2022

Related U.S. Application Data

(63) Continuation of application No. 16/862,936, filed on Apr. 30, 2020, now Pat. No. 11,302,080.

(60) Provisional application No. 62/843,855, filed on May 6, 2019.

(51) **Int. Cl.**

G06T 19/00	(2011.01)
G06F 3/0482	(2013.01)
G06F 3/04845	(2022.01)
G06N 3/08	(2023.01)
G06N 5/02	(2023.01)
G06N 20/00	(2019.01)

(52) U.S. Cl.

CPC **G06T 19/006** (2013.01); **G06F 3/0482**
(2013.01); **G06F 3/04845** (2013.01); **G06N**
3/08 (2013.01); **G06N 5/02** (2013.01); **G06N**
20/00 (2019.01)

(58) **Field of Classification Search**

CPC G06T 19/006; G06F 3/0482; G06F
3/04845; G06N 3/08; G06N 5/02; G06N
20/00

See application file for complete search history.

(56) **References Cited**

U.S. PATENT DOCUMENTS

10,055,681	B2 *	8/2018	Brown	G06N 3/008
10,103,964	B1	12/2019	Valgardsson et al.	
2003/0177187	A1	9/2003	Levine et al.	
2005/0228940	A1 *	10/2005	Blacquire	G11B 27/105
2017/0213473	A1 *	7/2017	Ribeira	G09B 5/10
2018/0114264	A1 *	4/2018	Rafii	G06T 15/50
2018/0120928	A1 *	5/2018	Hayashida	G06F 3/011
2018/0232204	A1	8/2018	Ghatage et al.	

(Continued)


Primary Examiner — Grace Q Li

(74) *Attorney, Agent, or Firm* — Fernando & Partners,
LLP

(57) **ABSTRACT**

In some implementations, a method includes obtaining an objective for a computer-generated reality (CGR) representation of an objective-effectuator. In some implementations, the objective is associated with a plurality of time frames. In some implementations, the method includes determining a plurality of candidate plans that satisfy the objective. In some implementations, the method includes selecting a first candidate plan of the plurality of candidate plans based on a selection criterion. In some implementations, the method includes effectuating the first candidate plan in order to satisfy the objective. In some implementations, the first candidate plan triggers the CGR representation of the objective-effectuator to perform a series of actions over the plurality of time frames associated with the objective.

23 Claims, 9 Drawing Sheets



(56)

References Cited

U.S. PATENT DOCUMENTS

2018/0260023	A1 *	9/2018	Perlin	G06F 1/1626
2018/0314484	A1 *	11/2018	Pahud	G06F 3/1423
2018/0336415	A1	11/2018	Anorga et al.	
2019/0321730	A1	10/2019	Shetty et al.	
2020/0306640	A1	10/2020	Kolen et al.	
2020/0312003	A1	10/2020	Borovikov et al.	
2021/0166459	A1 *	6/2021	Miller, IV	G06F 1/1637

* cited by examiner

100




Figure 1




Figure 2

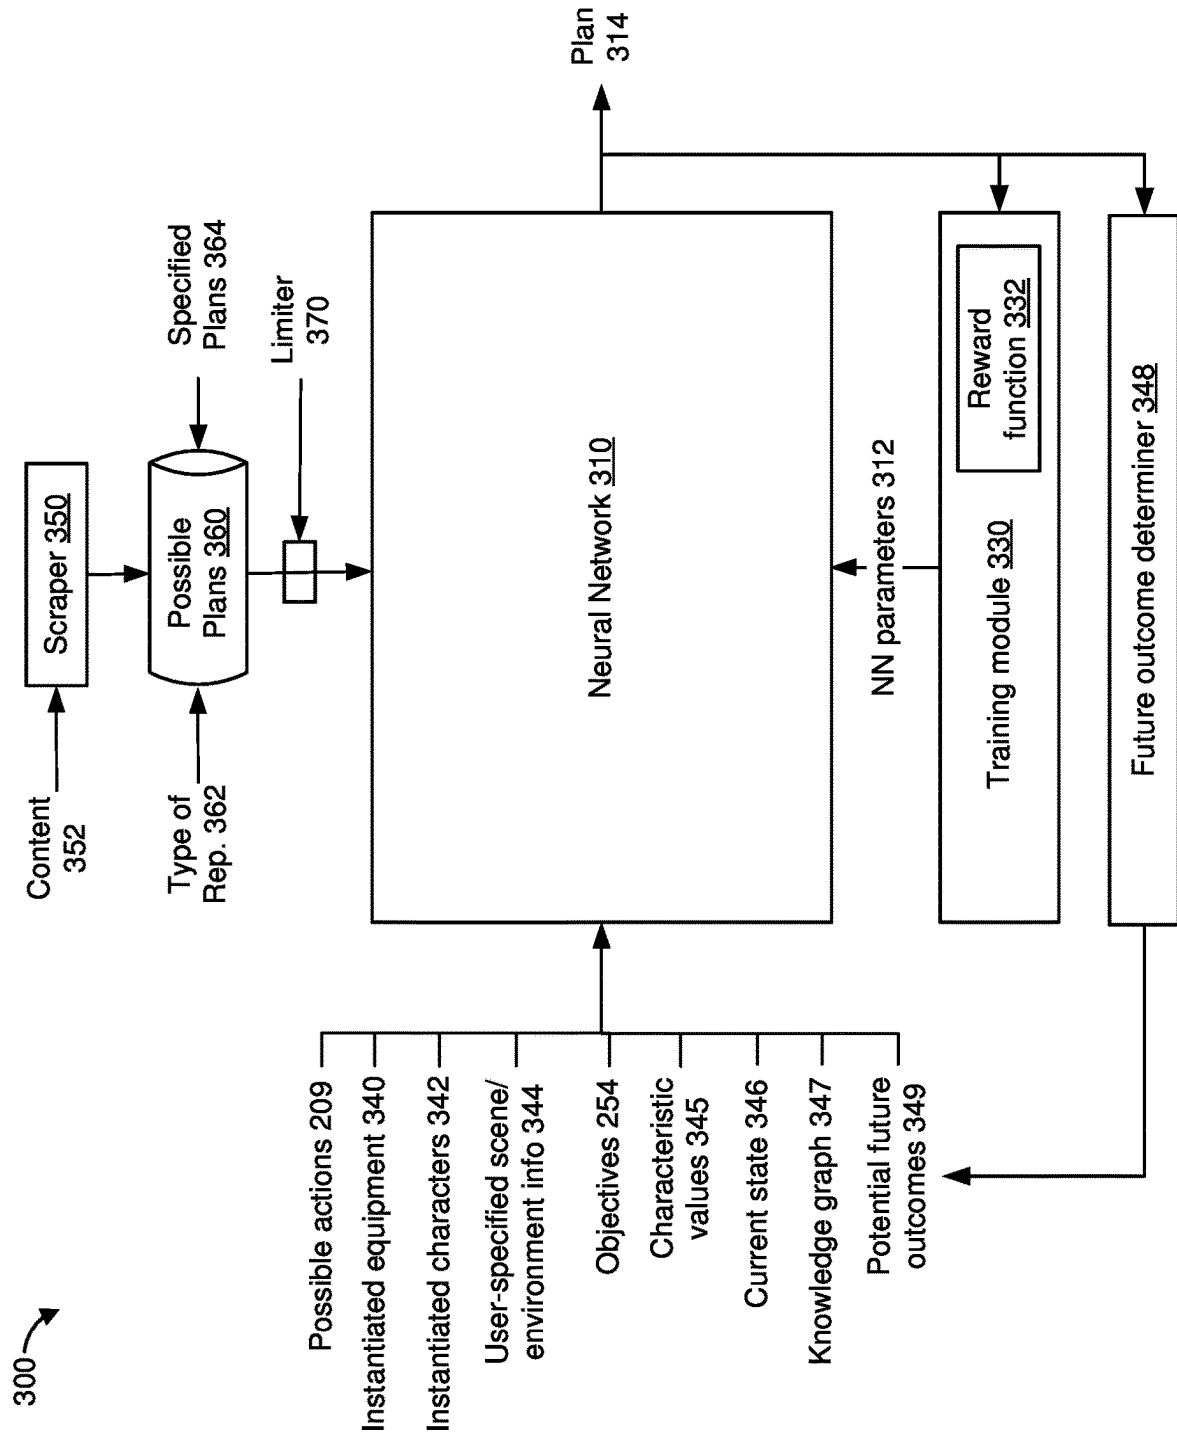


Figure 3A




Figure 3B

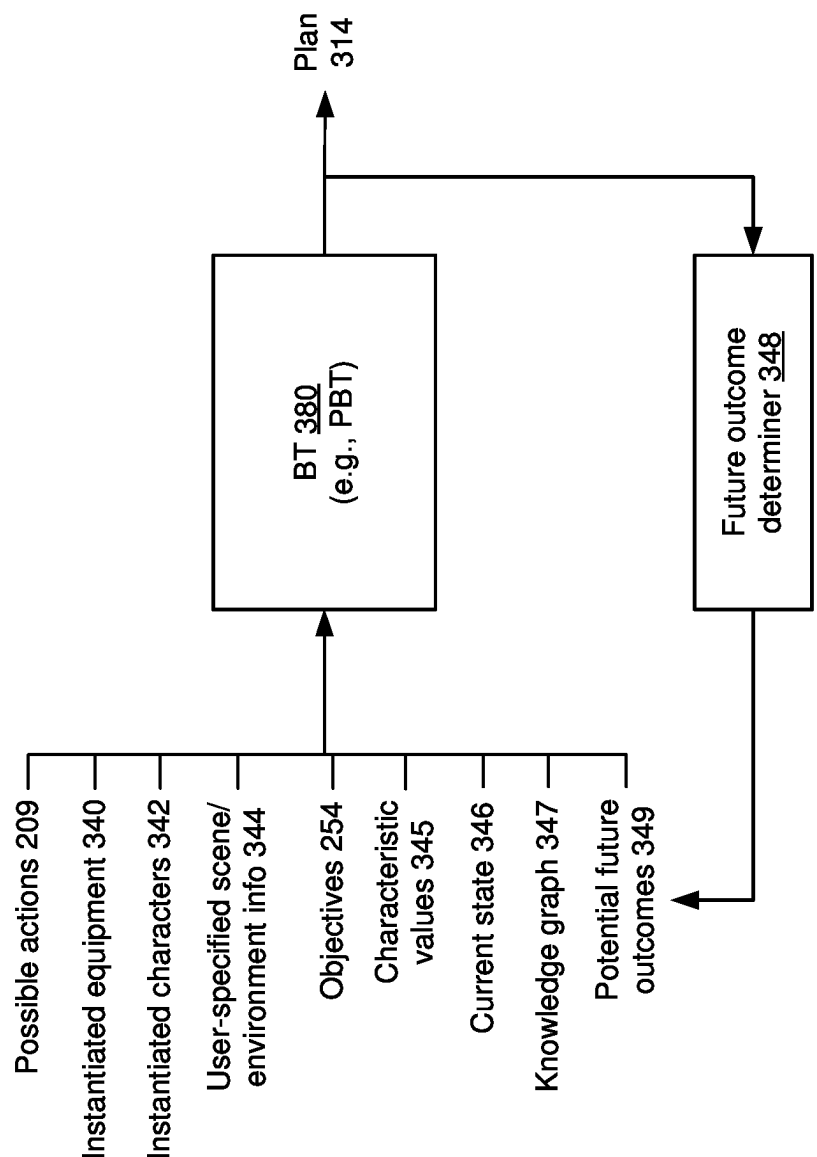
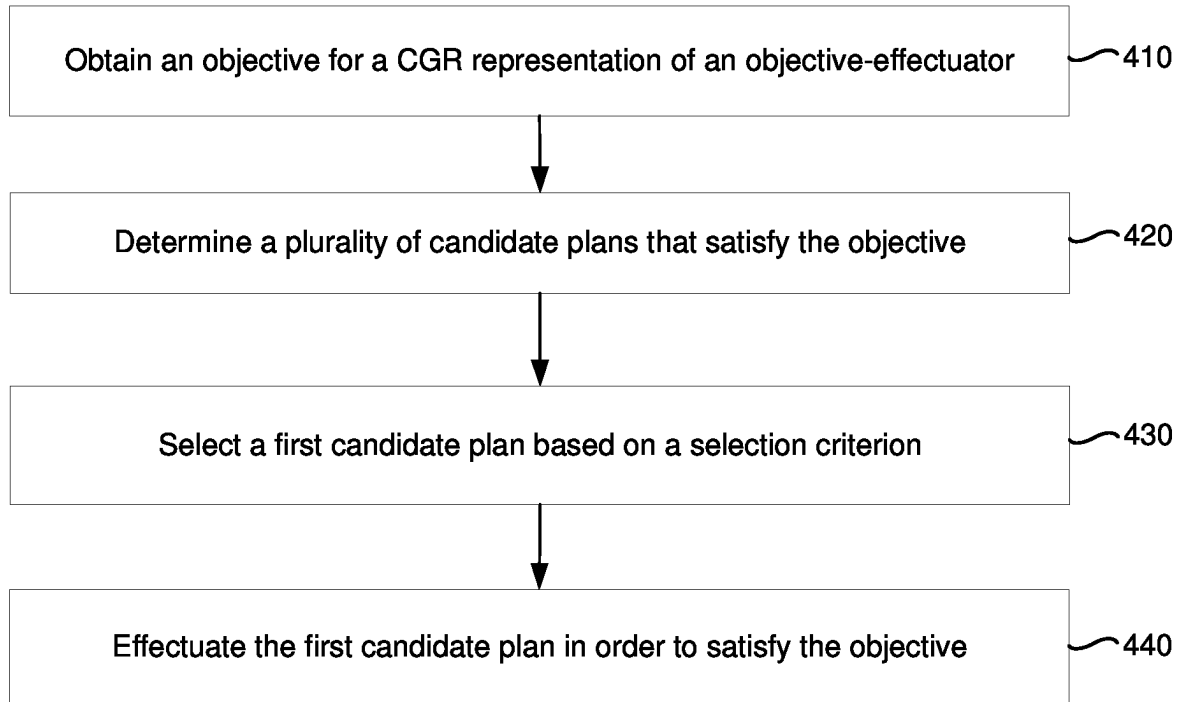
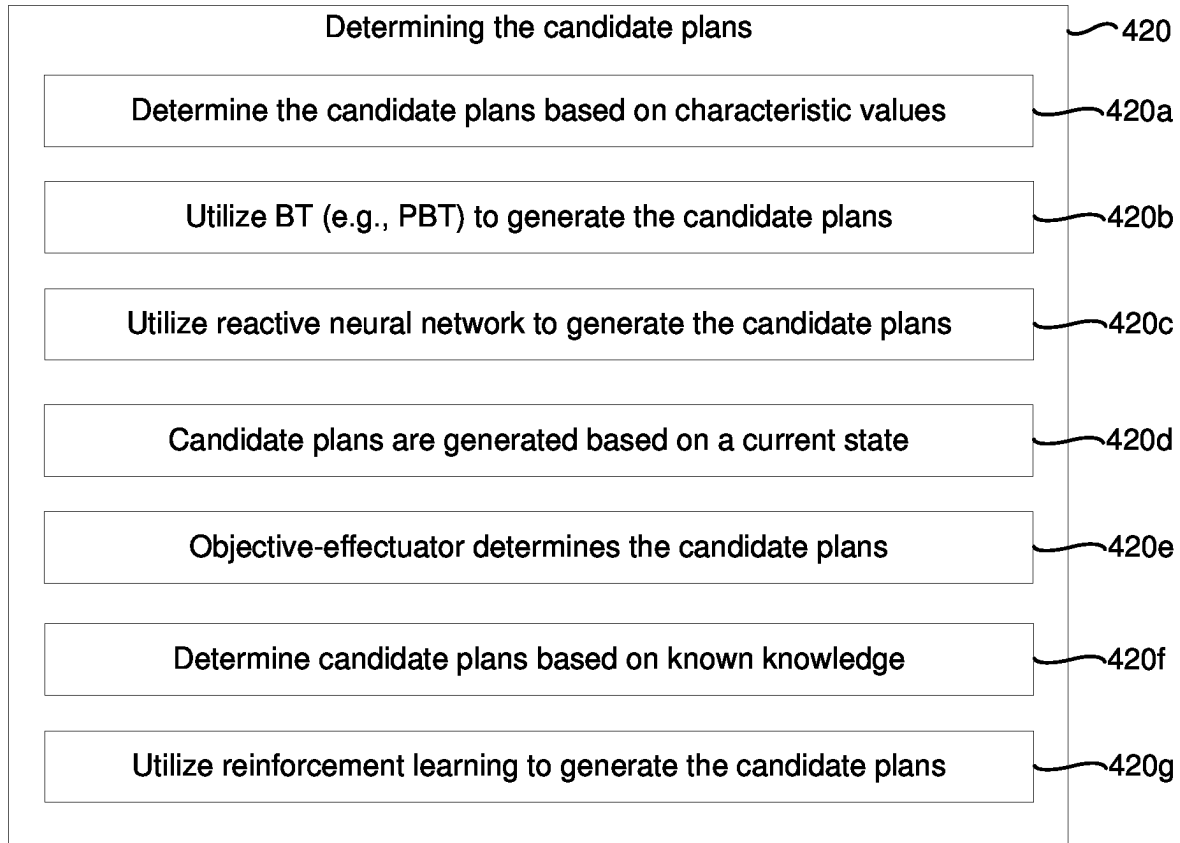
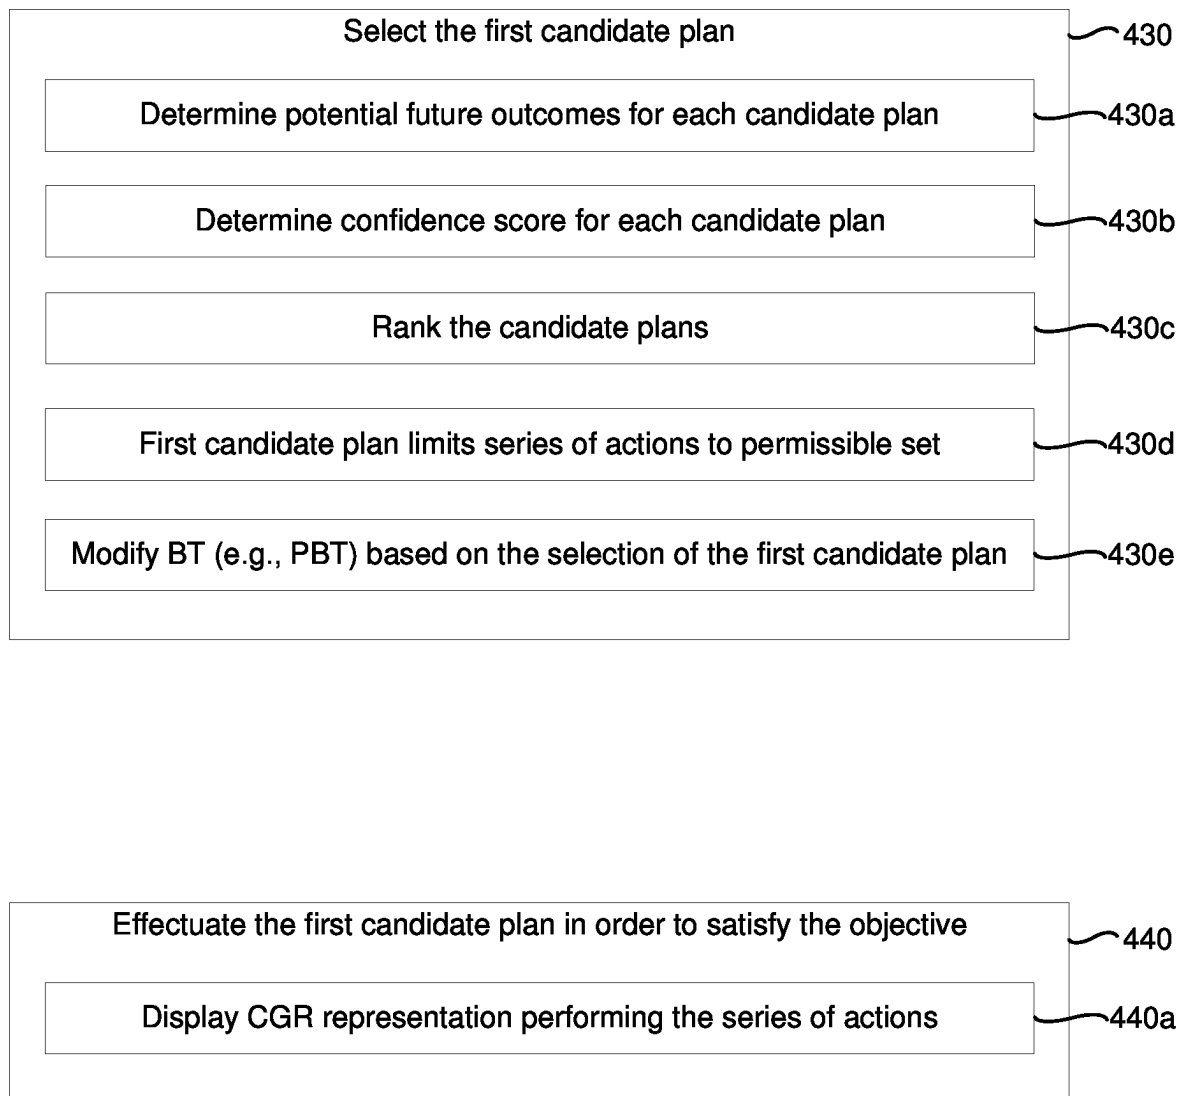
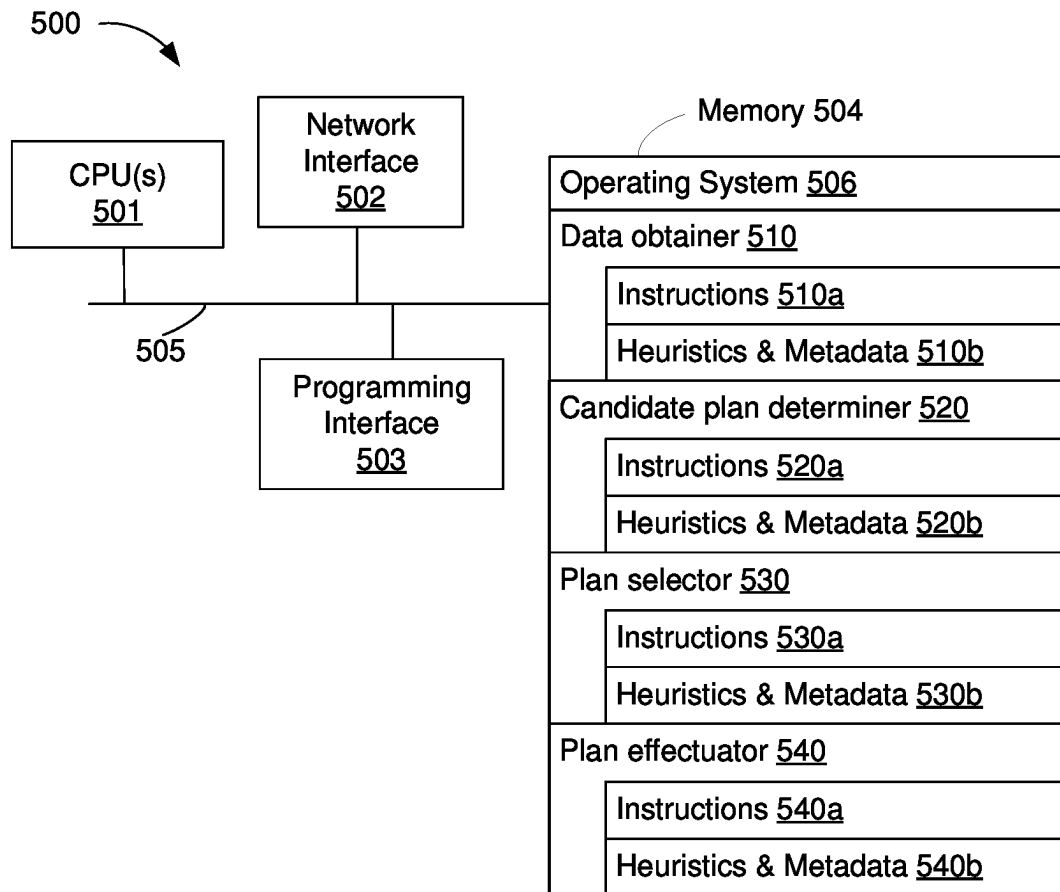


Figure 3C

400**Figure 4A**

**Figure 4B**

**Figure 4C**

**Figure 5**

1

PLANNER FOR AN OBJECTIVE-EFFECTUATOR

CROSS-REFERENCE TO RELATED APPLICATION

This application is a continuation application of and claims priority to U.S. patent application Ser. No. 16/862,936, filed on Apr. 30, 2020, which claims priority to U.S. Provisional Patent App. No. 62/843,855, filed on May 6, 2019, which are hereby incorporated by reference in their entirety.

TECHNICAL FIELD

The present disclosure generally relates to a planner for an objective-effectuator.

BACKGROUND

Some devices are capable of generating and presenting computer-generated reality (CGR) environments. Some CGR environments include virtual environments that are simulated replacements of physical environments. Some CGR environments include augmented environments that are modified versions of physical environments. Some devices that present CGR environments include mobile communication devices such as smartphones, head-mountable displays (HMDs), eyeglasses, heads-up displays (HUDs), and optical projection systems. Most previously available devices that present CGR environments are ineffective at presenting representations of certain objects. For example, some previously available devices that present CGR environments are unsuitable for presenting representations of objects that are associated with an action.

BRIEF DESCRIPTION OF THE DRAWINGS

So that the present disclosure can be understood by those of ordinary skill in the art, a more detailed description may be had by reference to aspects of some illustrative implementations, some of which are shown in the accompanying drawings.

FIG. 1 is a diagram of an example operating environment in accordance with some implementations.

FIG. 2 is a block diagram of an example system in accordance with some implementations.

FIGS. 3A-3C are block diagrams of example planners in accordance with some implementations.

FIGS. 4A-4C are flowchart representations of a method of generating a plan for an objective-effectuator in accordance with some implementations.

FIG. 5 is a block diagram of a device enabled with various components of the planner in accordance with some implementations.

In accordance with common practice the various features illustrated in the drawings may not be drawn to scale. Accordingly, the dimensions of the various features may be arbitrarily expanded or reduced for clarity. In addition, some of the drawings may not depict all of the components of a given system, method or device. Finally, like reference numerals may be used to denote like features throughout the specification and figures.

SUMMARY

Various implementations disclosed herein include devices, systems, and methods for determining a plan to

2

satisfy an objective of an objective-effectuator. In various implementations, a device includes a non-transitory memory and one or more processors coupled with the non-transitory memory. In some implementations, a method includes obtaining an objective for a computer-generated reality (CGR) representation of an objective-effectuator. In some implementations, the objective is associated with a plurality of time frames. In some implementations, the method includes determining a plurality of candidate plans that satisfy the objective. In some implementations, the method includes selecting a first candidate plan of the plurality of candidate plans based on a selection criterion. In some implementations, the method includes effectuating the first candidate plan in order to satisfy the objective. In some implementations, the first candidate plan triggers the CGR representation of the objective-effectuator to perform a series of actions over the plurality of time frames associated with the objective.

In accordance with some implementations, a device includes one or more processors, a non-transitory memory, and one or more programs. In some implementations, the one or more programs are stored in the non-transitory memory and are executed by the one or more processors. In some implementations, the one or more programs include instructions for performing or causing performance of any of the methods described herein. In accordance with some implementations, a non-transitory computer readable storage medium has stored therein instructions that, when executed by one or more processors of a device, cause the device to perform or cause performance of any of the methods described herein. In accordance with some implementations, a device includes one or more processors, a non-transitory memory, and means for performing or causing performance of any of the methods described herein.

DESCRIPTION

Numerous details are described in order to provide a thorough understanding of the example implementations shown in the drawings. However, the drawings merely show some example aspects of the present disclosure and are therefore not to be considered limiting. Those of ordinary skill in the art will appreciate that other effective aspects and/or variants do not include all of the specific details described herein. Moreover, well-known systems, methods, components, devices and circuits have not been described in exhaustive detail so as not to obscure more pertinent aspects of the example implementations described herein.

A physical environment refers to a physical world that people can sense and/or interact with without aid of electronic systems. Physical environments, such as a physical park, include physical articles, such as physical trees, physical buildings, and physical people. People can directly sense and/or interact with the physical environment, such as through sight, touch, hearing, taste, and smell.

In contrast, a computer-generated reality (CGR) environment refers to a wholly or partially simulated environment that people sense and/or interact with via an electronic system. In CGR, a subset of a person's physical motions, or representations thereof, are tracked, and, in response, one or more characteristics of one or more virtual objects simulated in the CGR environment are adjusted in a manner that comports with at least one law of physics. For example, a CGR system may detect a person's head turning and, in response, adjust graphical content and an acoustic field presented to the person in a manner similar to how such views and sounds would change in a physical environment.

In some situations (e.g., for accessibility reasons), adjustments to characteristic(s) of virtual object(s) in a CGR environment may be made in response to representations of physical motions (e.g., vocal commands).

A person may sense and/or interact with a CGR object using any one of their senses, including sight, sound, touch, taste, and smell. For example, a person may sense and/or interact with audio objects that create 3D or spatial audio environment that provides the perception of point audio sources in 3D space. In another example, audio objects may enable audio transparency, which selectively incorporates ambient sounds from the physical environment with or without computer-generated audio. In some CGR environments, a person may sense and/or interact only with audio objects.

Examples of CGR include virtual reality and mixed reality.

A virtual reality (VR) environment refers to a simulated environment that is designed to be based entirely on computer-generated sensory inputs for one or more senses. A VR environment comprises a plurality of virtual objects with which a person may sense and/or interact. For example, computer-generated imagery of trees, buildings, and avatars representing people are examples of virtual objects. A person may sense and/or interact with virtual objects in the VR environment through a simulation of the person's presence within the computer-generated environment, and/or through a simulation of a subset of the person's physical movements within the computer-generated environment.

In contrast to a VR environment, which is designed to be based entirely on computer-generated sensory inputs, a mixed reality (MR) environment refers to a simulated environment that is designed to incorporate sensory inputs from the physical environment, or a representation thereof, in addition to including computer-generated sensory inputs (e.g., virtual objects). On a virtuality continuum, a mixed reality environment is anywhere between, but not including, a wholly physical environment at one end and virtual reality environment at the other end.

In some MR environments, computer-generated sensory inputs may respond to changes in sensory inputs from the physical environment. Also, some electronic systems for presenting an MR environment may track location and/or orientation with respect to the physical environment to enable virtual objects to interact with real objects (that is, physical articles from the physical environment or representations thereof). For example, a system may account for movements so that a virtual tree appears stationery with respect to the physical ground.

Examples of mixed realities include augmented reality and augmented virtuality.

An augmented reality (AR) environment refers to a simulated environment in which one or more virtual objects are superimposed over a physical environment, or a representation thereof. For example, an electronic system for presenting an AR environment may have a transparent or translucent display through which a person may directly view the physical environment. The system may be configured to present virtual objects on the transparent or translucent display, so that a person, using the system, perceives the virtual objects superimposed over the physical environment. Alternatively, a system may have an opaque display and one or more imaging sensors that capture images or video of the physical environment, which are representations of the physical environment. The system composites the images or video with virtual objects, and presents the composition on the opaque display. A person, using the

system, indirectly views the physical environment by way of the images or video of the physical environment, and perceives the virtual objects superimposed over the physical environment. As used herein, a video of the physical environment shown on an opaque display is called "pass-through video," meaning a system uses one or more image sensor(s) to capture images of the physical environment, and uses those images in presenting the AR environment on the opaque display. Further alternatively, a system may have a projection system that projects virtual objects into the physical environment, for example, as a hologram or on a physical surface, so that a person, using the system, perceives the virtual objects superimposed over the physical environment.

An augmented reality environment also refers to a simulated environment in which a representation of a physical environment is transformed by computer-generated sensory information. For example, in providing pass-through video, a system may transform one or more sensor images to impose a select perspective (e.g., viewpoint) different than the perspective captured by the imaging sensors. As another example, a representation of a physical environment may be transformed by graphically modifying (e.g., enlarging) portions thereof, such that the modified portion may be representative but not photorealistic versions of the originally captured images. As a further example, a representation of a physical environment may be transformed by graphically eliminating or obfuscating portions thereof.

An augmented virtuality (AV) environment refers to a simulated environment in which a virtual or computer generated environment incorporates one or more sensory inputs from the physical environment. The sensory inputs may be representations of one or more characteristics of the physical environment. For example, an AV park may have virtual trees and virtual buildings, but people with faces photorealistically reproduced from images taken of physical people. As another example, a virtual object may adopt a shape or color of a physical article imaged by one or more imaging sensors. As a further example, a virtual object may adopt shadows consistent with the position of the sun in the physical environment.

There are many different types of electronic systems that enable a person to sense and/or interact with various CGR environments. Examples include head mounted systems, projection-based systems, heads-up displays (HUDs), vehicle windshields having integrated display capability, windows having integrated display capability, displays formed as lenses designed to be placed on a person's eyes (e.g., similar to contact lenses), headphones/earphones, speaker arrays, input systems (e.g., wearable or handheld controllers with or without haptic feedback), smartphones, tablets, and desktop/laptop computers. A head mounted system may have one or more speaker(s) and an integrated opaque display. Alternatively, a head mounted system may be configured to accept an external opaque display (e.g., a smartphone). The head mounted system may incorporate one or more imaging sensors to capture images or video of the physical environment, and/or one or more microphones to capture audio of the physical environment. Rather than an opaque display, a head mounted system may have a transparent or translucent display. The transparent or translucent display may have a medium through which light representative of images is directed to a person's eyes. The display may utilize digital light projection, OLEDs, LEDs, uLEDs, liquid crystal on silicon, laser scanning light source, or any combination of these technologies. The medium may be an optical waveguide, a hologram medium, an optical combiner, an optical reflector, or any combination thereof. In one

implementation, the transparent or translucent display may be configured to become opaque selectively. Projection-based systems may employ retinal projection technology that projects graphical images onto a person's retina. Projection systems also may be configured to project virtual objects into the physical environment, for example, as a hologram or on a physical surface.

The present disclosure provides methods, systems, and/or devices for generating a plan that satisfies an objective of an objective-effectuator. An objective-effectuator performs a set of one or more actions in order to effectuate an objective of the objective-effectuator. Some objectives span multiple time frames. For example, some objectives span several minutes, hours or days. A planner generates a plan for satisfying an objective that satisfies multiple time frames. The planner identifies various candidate plans that satisfy the objective. Each candidate plan is associated with a corresponding confidence score. The planner selects one of the candidate plans based on the respective confidence scores. The objective-effectuator effectuates the selected plan in order to satisfy the objective of the objective-effectuator. The selected plan triggers a computer-generated reality (CGR) representation of the objective-effectuator to perform a series of actions over the time frames associated with the objective. For example, if the objective spans twenty minutes, then the selected plan triggers the CGR representation of the objective-effectuator to perform various actions over the twenty minutes in order to satisfy the objective.

FIG. 1 is a block diagram of an example operating environment 100 in accordance with some implementations. While pertinent features are shown, those of ordinary skill in the art will appreciate from the present disclosure that various other features have not been illustrated for the sake of brevity and so as not to obscure more pertinent aspects of the example implementations disclosed herein. To that end, as a non-limiting example, the operating environment 100 includes a controller 102 and an electronic device 103. In the example of FIG. 1, the electronic device 103 is being held by a user 10. In some implementations, the electronic device 103 includes a smartphone, a tablet, a laptop, or the like.

As illustrated in FIG. 1, the electronic device 103 presents a computer-generated reality (CGR) environment 106. In some implementations, the CGR environment 106 is generated by the controller 102 and/or the electronic device 103. In some implementations, the CGR environment 106 includes a virtual environment that is a simulated replacement of a physical environment. In other words, in some implementations, the CGR environment 106 is synthesized by the controller 102 and/or the electronic device 103. In such implementations, the CGR environment 106 is different from the physical environment where the electronic device 103 is located. In some implementations, the CGR environment 106 includes an augmented environment that is a modified version of a physical environment. For example, in some implementations, the controller 102 and/or the electronic device 103 modify (e.g., augment) the physical environment where the electronic device 103 is located in order to generate the CGR environment 106. In some implementations, the controller 102 and/or the electronic device 103 generate the CGR environment 106 by simulating a replica of the physical environment where the electronic device 103 is located. In some implementations, the controller 102 and/or the electronic device 103 generate the CGR environment 106 by removing and/or adding items from the simulated replica of the physical environment where the electronic device 103 is located.

In some implementations, the CGR environment 106 includes various CGR representations of objective-effectuators, such as a boy action figure representation 108a that represents a boy objective-effectuator, a girl action figure representation 108b that represents a girl objective-effectuator, a robot representation 108c that represents a robot objective-effectuator, and a drone representation 108d that represents a drone objective-effectuator. In some implementations, the objective-effectuators represent (e.g., model behavior of) characters from fictional materials, such as movies, video games, comics, and novels. For example, the boy action figure representation 108a represents (e.g., models behavior of) a 'boy action figure' character from a fictional comic, and the girl action figure representation 108b represents (e.g., models behavior of) a 'girl action figure' character from a fictional video game. In some implementations, the CGR environment 106 includes objective-effectuators that represent (e.g., model behavior of) respective characters from different fictional materials (e.g., from different movies/games/comics/novels). In various implementations, the objective-effectuators model behavior of physical entities (e.g., tangible objects). For example, in some implementations, the objective-effectuators represent equipment (e.g., machinery such as planes, tanks, robots, cars, etc.). In the example of FIG. 1, the robot representation 108c represents a robot object-effectuator that models the behavior of a robot, and the drone representation 108d represents a drone objective-effectuator that models the behavior of a drone. In some implementations, the objective-effectuators represent (e.g., model behavior of) fictional entities (e.g., characters, equipment, etc.) from fictional materials. In some implementations, the objective-effectuators represent (e.g., model behavior of) entities from a physical environment, including entities located inside and/or outside of the CGR environment 106. In various implementations, an objective-effectuator models the behavior of an entity by manipulating a CGR representation of the objective-effectuator in order to provide an appearance that the CGR representation of the objective-effectuator is performing a set of one or more actions that are within a similarity threshold of actions that the entity performs.

In various implementations, the objective-effectuators perform one or more actions in order to effectuate (e.g., complete/satisfy/achieve) one or more objectives. In some implementations, the objective-effectuators perform a sequence of actions. In some implementations, the controller 102 and/or the electronic device 103 determine the actions that the objective-effectuators perform. In some implementations, the actions of the objective-effectuators are within a degree of similarity to (e.g., within a similarity threshold of) actions that the corresponding entities (e.g., characters or objects) perform in the fictional material. In the example of FIG. 1, the girl action figure representation 108b is performing the action of flying (e.g., because the corresponding 'girl action figure' character is capable of flying, and/or the 'girl action figure' character frequently flies in the fictional materials). In the example of FIG. 1, the drone representation 108d is performing the action of hovering (e.g., because drones in physical environments are capable of hovering). In some implementations, the controller 102 and/or the electronic device 103 obtain the actions for the objective-effectuators. For example, in some implementations, the controller 102 and/or the electronic device 103 receive the actions for the objective-effectuators from a remote server that determines (e.g., selects) the actions.

In various implementations, an objective-effectuator performs an action in order to satisfy (e.g., complete or achieve)

an objective. In some implementations, an objective-effectuator is associated with a particular objective, and the objective-effectuator performs actions that improve the likelihood of satisfying that particular objective. In some implementations, CGR representations of the objective-effectuators are referred to as CGR objects. In some implementations, an objective-effectuator representing (e.g., modeling the behavior of) a character is referred to as a character objective-effectuator. In some implementations, a character objective-effectuator performs actions to effectuate a character objective. In some implementations, an objective-effectuator representing (e.g., modeling the behavior of) an equipment is referred to as an equipment objective-effectuator. In some implementations, an equipment objective-effectuator performs actions to effectuate an equipment objective. In some implementations, an objective effectuator representing (e.g., modeling the behavior of) an environment is referred to as an environmental objective-effectuator. In some implementations, an environmental objective-effectuator performs environmental actions to effectuate an environmental objective.

In some implementations, the CGR environment **106** is generated based on a user input from the user **10**. For example, in some implementations, the electronic device **103** receives a user input indicating a terrain for the CGR environment **106**. In such implementations, the controller **102** and/or the electronic device **103** configure the CGR environment **106** such that the CGR environment **106** includes the terrain indicated via the user input. In some implementations, the user input indicates environmental conditions for the CGR environment **106**. In such implementations, the controller **102** and/or the electronic device **103** configure the CGR environment **106** to have the environmental conditions indicated by the user input. In some implementations, the environmental conditions include one or more of temperature, humidity, pressure, visibility, ambient light level, ambient sound level, time of day (e.g., morning, afternoon, evening, or night), and precipitation (e.g., overcast, rain, or snow). In some implementations, the user input specifies a time period for the CGR environment **106**. In such implementations, the controller **102** and/or the electronic device **103** maintain and present the CGR environment **106** during the specified time period.

In some implementations, the controller **102** and/or the electronic device **103** determine (e.g., generate) actions for the objective-effectuators based on a user input from the user **10**. For example, in some implementations, the electronic device **103** receives a user input indicating placement of the CGR representations of the objective-effectuators. In such implementations, the controller **102** and/or the electronic device **103** position the CGR representations of the objective-effectuators in accordance with the placement indicated by the user input. In some implementations, the user input indicates specific actions that the objective-effectuators are permitted to perform. In such implementations, the controller **102** and/or the electronic device **103** select the actions for the objective-effectuators from the specific actions indicated by the user input. In some implementations, the controller **102** and/or the electronic device **103** forgo actions that are not among the specific actions indicated by the user input.

In various implementations, the controller **102** and/or the electronic device **103** generate a plan to satisfy an objective of an objective-effectuator. In some implementations, the plan triggers the CGR representation of the objective-effectuator to perform a series of actions over a set of time frames associated with the objective. In some implementations, the plan specifies types of actions for the CGR representation of

the objective-effectuator to perform in order to advance (e.g., satisfy) the objective of the objective-effectuator. In some implementations, the plan includes specific actions for the CGR representation of the objective-effectuator to perform in order to advance the objective of the objective-effectuator.

In some implementations, the plan includes a timeline which orders the actions for the CGR representation of the objective-effectuator over the set of time frames associated with the objective. In some implementations, the plan distributes the actions for the CGR representation of the objective-effectuator across the set of time frames associated with the objective. For example, in some implementations, the plan distributes the actions across the set of time frames such that the actions are spread approximately evenly across the set of time frames. In some implementations, the plan includes performing at least one action during each of the time frame in the set of time frames.

In some implementations, the controller **102** and/or the electronic device **103** generate the plan for an objective-effectuator based on a plot template associated with the CGR environment **106**. In some implementations, the plot template defines a type of plot for the CGR environment **106**. Example plot templates include a mystery plot template, a disaster plot template, a comedy plot template and a rescue plot template. In some implementations, the plan triggers actions that are consistent with the plot template for the CGR environment **106**. For example, if the plot template for the CGR environment **106** is a comedy plot template, then the plan for an objective-effectuator triggers actions that satisfy a comedic threshold.

In some implementations, the user **10** wears a head-mountable device (HMD). In various implementations, the HMD operates in substantially the same manner as the electronic device **103** shown in FIG. 1. In some implementations, the HMD performs substantially the same operations as the electronic device **103** shown in FIG. 1. For example, in some implementations, the HMD, being worn by the user **10**, presents (e.g., displays) the CGR environment **106** according to various implementations. In some implementations, the HMD includes an integrated display (e.g., a built-in display) that displays the CGR environment **106**. In some implementations, the HMD includes a head-mountable enclosure. In various implementations, the head-mountable enclosure includes an attachment region to which another device with a display can be attached. For example, in some implementations, the electronic device **103** can be attached to the head-mountable enclosure. In various implementations, the head-mountable enclosure is shaped to form a receptacle for receiving another device that includes a display (e.g., the electronic device **103**). For example, in some implementations, the electronic device **103** slides/snaps into or otherwise attaches to the head-mountable enclosure. In some implementations, the display of the device attached to the head-mountable enclosure presents (e.g., displays) the CGR environment **106**.

FIG. 2 is a block diagram of an example system **200** that generates plans for various objective-effectuators in a CGR environment. To that end, the system includes planners **204a** . . . **204e**, objective-effectuator engines, and an emergent content engine **250**. In various implementations, the emergent content engine **250** provides objectives **254** for various objective-effectuators. The planners **204a** . . . **204e** generate plans **206a** . . . **206e** for the objective-effectuators based on the objectives **254**. The planners **204a** . . . **204e** provide the plans **206a** . . . **206e** to the objective-effectuator engines. The objective-effectuator engines generate actions

210a . . . 210e based on the plans **206a . . . 206e**. The actions **210a . . . 210e** are performed by CGR representations of the objective-effectuators.

In various implementations, the planners **204a . . . 204e** and objective-effectuator engines collectively generate the plans **206a . . . 206e** and the actions **210a . . . 210e**. In some implementations, the objective-effectuator engines generate the actions **210a . . . 210e** in accordance with the plans **206a . . . 206e**. In various implementations, the planners **204a . . . 204e** generate the plans **206a . . . 206e** based on the actions **210a . . . 210e**. In some implementations, the planners **204a . . . 204e** generate/adjust the plans **206a . . . 206e** based on historical/future actions.

In some implementations, a planner generates a plan that triggers a CGR representation of a particular objective-effectuator to perform actions in accordance with the plan. In the example of FIGS. 1 and 2, the planner **204a** generates the plan **206a** which triggers the boy action figure representation **108a** to perform the actions **210a**. The planner **204b** generates the plan **206b** which triggers the girl action figure representation **108b** to perform the actions **210b**. The planner **204c** generates the plan **206c** which triggers the robot representation **108c** to perform the actions **210c**. The planner **204d** generates the plan **206d** which triggers the drone representation **108d** to perform the actions **210d**. The planner **204e** generates the plan **206e** which triggers an environment of the CGR environment **106** to perform the actions **210e**.

In various implementations, the planners **204a . . . 204e** select the plans **206a . . . 206e** from a set of possible plans **206** stored in a datastore. In some implementations, the possible plans **206** are obtained from corresponding fictional source material. For example, in some implementations, the controller **102** and/or the electronic device **103** shown in FIG. 1 perform semantic analysis on actions performed in the fictional source material and determine (e.g., construct) the possible plans **206** based on the actions. In some implementations, the possible plans **206** vary for different objective-effectuators. For example, in some implementations, the possible plans **206** for the girl action figure representation **108b** include planning a rescue operation and planning a fight, whereas the possible plans **206** for the boy action figure representation **108a** include planning a comedy routine and planning a social gathering.

In some implementations, the planners **204a . . . 204e** generate the plans **206a . . . 206e** based on a set of possible actions **209** stored in a datastore. In some implementations, the set of possible actions **209** represent actions that the CGR representations of objective-effectuators are capable of performing in a CGR environment. For example, the set of possible actions **209** represent actions that the boy action figure representation **108a**, the girl action figure representation **108b**, the robot representation **108c** and/or the drone representation **108d** are capable of performing. In some implementations, the planners **204a . . . 204e** generate the plans **206a . . . 206e** such that the plans **206a . . . 206e** can be satisfied (e.g., carried out) with the set of possible actions **209**.

In the example of FIG. 2, the system **200** includes objective-effectuator engines that generate actions in accordance with corresponding plans. For example, a character engine **208a** generates the actions **210a** for the boy action figure representation **108a** in accordance with the plan **206a**. A character engine **208b** generates the actions **210b** for the girl action figure representation **108b** in accordance with the plan **206b**. An equipment engine **208c** generates the actions **210c** for the robot representation **108c** (e.g., responses for

the robot representation **108c**) in accordance with the plan **206c**. An equipment engine **208d** generates the actions **210d** for the drone representation **108d** (e.g., responses for the drone representation **108d**) in accordance with the plan **206d**. An environmental engine **208e** generates environmental responses **210e** for an environment of the CGR environment **106** based on the plan **206e**.

In various implementations, the planners **204a . . . 204e** generate the plans **206a . . . 206e** based on a set of parameters **211**. In some implementations, the parameters **211** include a set of initialization conditions for a CGR environment. In some implementations, the parameters **211** are provided by an operator (e.g., a human operator, for example, a user of a device). In some implementations, the parameters **211** are generated by a CGR environment. In some implementations, the parameters **211** are provided by objective-effectuators that are instantiated in the CGR environment. In some implementations, the parameters **211** include initial/end states **212** for the CGR environment **106**. In some implementations, the parameters **211** include environmental conditions **214** for the CGR environment **106**.

In some implementations, the planners **204a . . . 204e** generate the plans **206a . . . 206e** based on the initial/end states **212** for the CGR environment **106**. In some implementations, the CGR environment **106** is associated with a time duration (e.g., a few minutes, hours or days). For example, the CGR environment **106** is scheduled to last for the time duration associated with the CGR environment **106**. In some implementations, the initial/end states **212** indicate initial/end placements of the characters/equipment at the beginning of the time duration and/or at the end of the time duration associated with the CGR environment **106**. In such implementations, the plans **206a . . . 206e** are a function of the initial/end placements of the characters/equipment at the beginning/end of the time duration associated with the CGR environment **106**.

In some implementations, the planners **204a . . . 204e** generate the plans **206a . . . 206e** based on the environmental conditions **214** for the CGR environment **106**. In some implementations, the environmental conditions **214** specify weather conditions within the CGR environment **106**. For example, the environmental conditions **214** specify an amount of precipitation, an amount of snowfall and/or a visibility level in the CGR environment **106**. In some implementations, the environmental conditions **214** are specified by the user. In some implementations, the planners **204a . . . 204e** forgo plans that are adversely impacted by the environmental conditions **214** (e.g., by assigning a low confidence score to adversely-affected plans). For example, if the environmental conditions **214** indicate that it is raining in the CGR environment **106**, then a plan to light a camp fire in the CGR environment **106** is assigned a low confidence score. In various implementations, the environmental conditions **214** limit the plans **206a . . . 206e** generated by the planners **204a . . . 204e**.

In some implementations, the emergent content engine **250** generates the objectives **254** based on a set of possible objectives **252** that are stored in a datastore. In some implementations, the possible objectives **252** are obtained from corresponding fictional source material. For example, in some implementations, the possible objectives **252** for the girl action figure representation **108b** include saving lives, rescuing pets, and/or fighting crime. In some implementations, the planners **204a . . . 204e** utilize the objectives **254** to generate additional plans and/or to modify previously-generated plans. For example, in some implementations, the planner **204a** utilizes the objectives **254** to generate an

additional plan **206a** and/or modify a previously-generated plan **206a** for the boy action figure representation **108a**.

In various implementations, the character/equipment/environmental engines provide the actions **210a** . . . **210e** to a display engine **260** (e.g., a rendering and display pipeline). In some implementations, the display engine **260** modifies the CGR representations of the objective-effectuators and/or the environment of the CGR environment **106** based on the actions **210a** . . . **210e**. In various implementations, the display engine **260** modifies the CGR representations of the objective-effectuators such that the CGR representations of the objective-effectuator can be seen as performing the actions **210a** . . . **210e**. For example, if an action for the girl action figure representation **108b** is to fly, the display engine **260** moves the girl action figure representation **108b** within the CGR environment **106** in order to give the appearance that the girl action figure representation **108b** is flying within the CGR environment **106**.

FIG. 3A is a block diagram of an example planner **300** in accordance with some implementations. In some implementations, the planner **300** generates a plan **314** for a corresponding objective-effectuator (e.g., a character objective-effectuator, an equipment objective-effectuator or an environmental objective-effectuator). In some implementations, the plan **314** triggers a CGR representation of the objective-effectuator (e.g., the boy action figure representation **108a**, the girl action figure representation **108b**, the robot representation **108c**, and/or the drone representation **108d**) to perform a series of actions in order to satisfy an objective.

In some implementations, different instances of the planner **300** implement the planners **204a** . . . **204e**. For example, a first instance of the planner **300** implements the planner **204a** which generates the plan **206a** for the boy action figure representation **108a**. A second instance of the planner **300** implements the planner **204b** which generates the plans **206b** for the girl action figure representation **108b**. A third instance of the planner **300** implements the planner **204c** which generates the plan **206c** for the robot representation **108c**. A fourth instance of the planner **300** implements the planner **204d** which generates the plan **206d** for the drone representation **108d**. A fifth instance of the planner **300** implements the planner **204e** which generates the plan **206e** for the environment of the CGR environment **106**.

In various implementations, the planner **300** includes a neural network system **310** ("neural network **310**", hereinafter for the sake of brevity), a neural network training system **330** ("a training module **330**", hereinafter for the sake of brevity) that trains (e.g., configures) the neural network **310**, and a scraper **350** that provides possible plans **360** to the neural network **310**. In various implementations, the neural network **310** generates the plan **314** for a corresponding objective-effectuator based on various inputs.

In some implementations, the neural network **310** includes a long short-term memory (LSTM) recurrent neural network (RNN). In various implementations, the neural network **310** generates the plan **314** based on a function of the possible plans **360**. For example, in some implementations, the neural network **310** generates the plan **314** by selecting one of the possible plans **360**. In some implementations, the neural network **310** generates the plan **314** such that the plan **314** is within a degree of similarity to at least one of the possible plans **360**.

In some implementations, the neural network **310** generates the plan **314** based on the set of possible actions **209**. As described herein, in some implementations, the set of possible actions **209** include actions that a CGR representation

of an objective-effectuator is capable of performing in a CGR environment. In some implementations, the neural network **310** generates the plan **314** such that the plan **314** can be satisfied (e.g., carried out) based on the set of possible actions **209**.

In some implementations, the neural network **310** generates the plan **314** based on instantiated equipment representations **340**. In some implementations, the instantiated equipment representations **340** refer to equipment objective-effectuators that are instantiated in the CGR environment. For example, referring to FIG. 1, the instantiated equipment representations **340** include the robot representation **108c** and the drone representation **108d** in the CGR environment **106**. In some implementations, the plan **314** includes engaging with one or more of the instantiated equipment representations **340**. For example, referring to FIG. 1, in some implementations, the plan **206b** for the girl action figure representation **108b** includes collaborating with the drone representation **108d** to perform a search and rescue operation.

In some implementations, the neural network **310** generates the plan **314** based on instantiated character representations **342**. In some implementations, the instantiated character representations **342** refer to character objective-effectuators that are instantiated in the CGR environment. For example, referring to FIG. 1, the instantiated character representations **342** include the boy action figure representation **108a** and the girl action figure representation **108b** in the CGR environment **106**. In some implementations, the plan **314** includes engaging with one or more of the instantiated character representations **342**. For example, referring to FIG. 1, in some implementations, the plan **206a** for the boy action figure representation **108a** includes catching the girl action figure representation **108b**. In some implementations, the plan **206c** for the robot representation **108c** includes assisting the boy action figure representation **108a** in catching the girl action figure representation **108b**.

In some implementations, the neural network **310** generates the plan **314** based on user-specified scene/environment information **344**. In some implementations, the user-specified scene/environment information **344** includes the initial/end states **212** and/or the environmental conditions **214** shown in FIG. 2. In some implementations, the plan **314** is a function of the initial/end states **212**. In some implementations, the initial/end states **212** indicate initial/end placement for CGR representation of the objective-effectuator, and the plan **314** includes performing operations that trigger the CGR representation of the objective-effectuator to move away/towards from/to the initial/end placement. In some implementations, the neural network **310** adjusts the plan **314** so that the plan **314** is better suited for the user-specified scene/environment information **344**. For example, the plan **314** includes a flying operation for the girl action figure representation **108b** when the user-specified scene/environment information **344** indicate that the skies within the CGR environment **106** are clear. In some implementations, the neural network **310** forgoes the plan **314** if the plan **314** is unsuitable for the environment indicated by the user-specified scene/environment information **344**. For example, if the plan **314** includes flying the drone representation **108d**, then the neural network **310** forgoes the plan **314** when the user-specified scene/environment information **344** indicate high winds within the CGR environment **106**.

In various implementations, the neural network **310** generates the plan **314** based on the objectives **254** from the emergent content engine **250**. In some implementations, the neural network **310** generates the plan **314** in order to satisfy

the objective 254 from the emergent content engine 250. In some implementations, the neural network 310 evaluates the possible plans 360 with respect to the objectives 254. In such implementations, the neural network 310 generates the plan 314 by selecting one of the possible plans 360 that satisfies the objectives 254 and forgoes selection of the possible plans 360 that do not satisfy the objectives 254.

In various implementations, the neural network 310 generates the plan 314 based on one or more characteristic values 345 associated with the objective-effectuator. In some implementations, the one or more characteristic values 345 indicate one or more physical characteristics (e.g., structural characteristics) of the CGR representation of the objective-effectuator. For example, the one or more characteristic values 345 indicate a body material of the CGR representation of the objective-effectuator. In such implementations, the plan 314 utilizes the physical characteristics that the CGR representation possesses, and does not utilize the physical characteristics that the CGR representation does not possess. For example, if the CGR representation is made from wax, then the plan 314 avoids going into an area of the CGR environment that is designated as a hot area where there is a risk of melting the wax.

In some implementations, the one or more characteristic values 345 indicate an accessory that the CGR representation of the objective-effectuator has (e.g., a jet pack for flying). In such implementations, the plan 314 utilizes the accessories that the CGR representation has, and avoids accessories that the CGR representation does not have. For example, if the CGR representation has the jet pack accessory, then the plan 314 may include flying. However, if the CGR representation does not have the jet pack accessory, then the plan 314 may not include flying or the plan 314 may include taking a CGR airplane to fly.

In some implementations, the one or more characteristic values indicate one or more behavioral characteristics of the CGR representation of the objective-effectuator. In some implementations, the behavioral characteristics include long-term personality traits such as a level of aggressiveness, a level of patience, a level of politeness, etc. In some implementations, the behavioral characteristics include short-term behavioral attributes such as a mood of the CGR representation of the objective-effectuator. In some implementations, the plan 314 includes actions that rely on behavioral traits which the CGR representation possesses. For example, if the CGR representation has a relatively high level of aggressiveness, then the plan 314 includes fighting. In some implementations, the plan 314 forgoes actions that rely on behavioral traits which the CGR representation does not possess. For example, if the CGR representation has a relatively low level of patience, then the plan 314 does not include waiting.

In some implementations, the neural network 310 obtains a current state 346 of the CGR representation of the objective-effectuator, and generates the plan 314 based on the current state 346 of the CGR representation of the objective-effectuator. In some implementations, the neural network 310 obtains the current state 346 from an emphatic engine that evaluates the current state 346 of the CGR representation of the objective-effectuator. In some implementations, the emphatic engine determines the current state 346 based on a body pose of the CGR representation. In some implementations, the emphatic engine determines the current state 346 based on a mood of the CGR representation of the objective-effectuator. In some implementations, the emphatic engine determines the current state 346 based on actions that the CGR representation is currently performing

and/or has recently performed (e.g., within a threshold time prior to a current time). In some implementations, the emphatic engine determines the current state 346 based on an objective of the CGR representation of the objective-effectuator.

In some implementations, the neural network 310 obtains a portion of a knowledge graph 347 associated with the CGR environment. In some implementations, the knowledge graph 347 includes state information that represents one or more states (e.g., a current state and/or historical states) of the CGR environment. In some implementations, the portion of the knowledge graph 347 represents knowledge of the objective-effectuator. In some implementations, the neural network 310 generates the plan 314 based on the portion of the knowledge graph 347. As such, in some implementations, the plan 314 is generated based on the knowledge of the objective-effectuator.

In some implementations, a future outcome determiner 348 determines potential future outcomes 349 for the plan 314, and the neural network 310 adjusts the plan 314 based on the potential future outcomes 349. In some implementations, the future outcome determiner 348 determines the potential future outcomes 349 for various candidate plans (e.g., the candidate plans 314a . . . 314n shown in FIG. 3B), and the neural network 310 selects the plan 314 based on the potential future outcomes 349 for the candidate plans. In some implementations, the neural network 310 selects the candidate plan with the most favorable potential future outcome 349. In some implementations, the neural network 310 selects the candidate plan with the potential future outcome 349 in which the objective 254 for the objective-effectuator is satisfied. In some implementations, the future outcome determiner 348 determines the potential future outcomes 349 by simulating the plan 314 (e.g., by executing the plan 314 in a test CGR environment).

In various implementations, the training module 330 trains the neural network 310. In some implementations, the training module 330 provides neural network (NN) parameters 312 to the neural network 310. In some implementations, the neural network 310 includes a model of neurons, and the neural network parameters 312 represent weights for the neurons. In some implementations, the training module 330 generates (e.g., initializes/initiates) the neural network parameters 312, and refines the neural network parameters 312 based on the plan 314 generated by the neural network 310.

In some implementations, the training module 330 includes a reward function 332 that utilizes reinforcement learning to train the neural network 310. In some implementations, the reward function 332 assigns a positive reward to plans that are desirable, and a negative reward to plans that are undesirable. In some implementations, during a training phase, the training module 330 compares the plans with verification data that includes verified plans. In such implementations, if the plan 314 is within a degree of similarity to the verified plans, then the training module 330 stops training the neural network 310. However, if the plan 314 is not within the degree of similarity to the verified plans, then the training module 330 continues to train the neural network 310. In various implementations, the training module 330 updates the neural network parameters 312 during/after the training.

In various implementations, the scraper 350 scrapes content 352 to identify the possible plans 360. In some implementations, the content 352 includes movies, video games, comics, novels, and fan-created content such as blogs and commentary. In some implementations, the scraper 350

utilizes various methods, systems, and devices associated with content scraping to scrape the content 352. For example, in some implementations, the scraper 350 utilizes one or more of text pattern matching, HTML (Hyper Text Markup Language) parsing, DOM (Document Object Model) parsing, image processing, and audio analysis in order to scrape the content 352 and identify the possible plans 360. In some implementations, the scraper 350 extracts actions from the content 352 and performs semantic analysis on the extracted actions to generate the possible plans 360.

In some implementations, an objective-effectuator is associated with a type of representation 362, and the neural network 310 generates the plan 314 based on the type of representation 362 associated with the objective-effectuator. In some implementations, the type of representation 362 indicates physical characteristics of the objective-effectuator, such as characteristics relating to its appearance and/or feel (e.g., color, material type, texture, etc.). In some implementations, the neural network 310 generates the plan 314 based on the physical characteristics of the objective-effectuator. In some implementations, the type of representation 362 indicates behavioral characteristics of the objective-effectuator (e.g., aggressiveness, friendliness, etc.). In some implementations, the neural network 310 generates the plan 314 based on the behavioral characteristics of the objective-effectuator. For example, the neural network 310 generates a flight plan for the boy action figure representation 108a in response to the behavioral characteristics including aggressiveness. In some implementations, the type of representation 362 indicates functional characteristics of the objective-effectuator (e.g., strength, speed, flexibility, etc.). In some implementations, the neural network 310 generates the plan 314 based on the functional characteristics of the objective-effectuator. For example, the neural network 310 generates a flight plan for the girl action figure representation 108b in response to the functional characteristics including flying. In some implementations, the type of representation 362 is determined based on a user input. In some implementations, the type of representation 362 is determined based on a combination of rules.

In some implementations, the neural network 310 generates the plan 314 based on specified plans 364. In some implementations, the specified plans 364 are provided by an entity that controls the fictional materials from where the character/equipment originated. For example, in some implementations, the specified plans 364 are provided (e.g., conceived of) by a movie producer, a video game creator, a novelist, etc. In some implementations, the possible plans 360 include the specified plans 364. As such, in some implementations, the neural network 310 generates the plan 314 by selecting a portion of the specified plans 364.

In some implementations, the possible plans 360 for an objective-effectuator are limited by a limiter 370. In some implementations, the limiter 370 restricts the neural network 310 from selecting a portion of the possible plans 360. In some implementations, the limiter 370 is controlled by the entity that controls (e.g., owns) the fictional materials from where the character/equipment originated. For example, in some implementations, the limiter 370 is controlled (e.g., operated and/or managed) by a movie producer, a video game creator, a novelist, etc. In some implementations, the limiter 370 and the neural network 310 are controlled/operated by different entities. In some implementations, the limiter 370 restricts the neural network 310 from generating plans that breach a criterion defined by the entity that controls the fictional materials.

FIG. 3B is a block diagram of the neural network 310 in accordance with some implementations. In the example of FIG. 3B, the neural network 310 includes an input layer 320, a first hidden layer 322, a second hidden layer 324, a classification layer 326, and a plan selector 328. While the neural network 310 includes two hidden layers as an example, those of ordinary skill in the art will appreciate from the present disclosure that one or more additional hidden layers are also present in various implementations. Adding additional hidden layers adds to the computational complexity and memory demands, but may improve performance for some applications.

In various implementations, the input layer 320 is coupled (e.g., configured) to receive various inputs. In the example of FIG. 3B, the input layer 320 receives inputs indicating the set of possible actions 209, the instantiated equipment representations 340, the instantiated character representations 342, the user-specified scene/environment information 344, the objectives 254 from the emergent content engine 250, the one or more characteristic values 345, the current state 346, and the portion of the knowledge graph 347. In some implementations, the neural network 310 includes a feature extraction module (not shown) that generates a feature stream (e.g., a feature vector) based on the set of possible actions 209, the instantiated equipment representations 340, the instantiated character representations 342, the user-specified scene/environment information 344, the objectives 254, the one or more characteristic values 345, the current state 346 and/or the portion of the knowledge graph 347. In such implementations, the feature extraction module provides the feature stream to the input layer 320. As such, in some implementations, the input layer 320 receives a feature stream that is a function of the set of possible actions 209, the instantiated equipment representations 340, the instantiated character representations 342, the user-specified scene/environment information 344, the objectives 254, the one or more characteristic values 345, the current state 346 and/or the portion of the knowledge graph 347. In various implementations, the input layer 320 includes a number of LSTM logic units 320a, which are also referred to as model(s) of neurons by those of ordinary skill in the art. In some such implementations, an input matrix from the features to the LSTM logic units 320a includes rectangular matrices. The size of this matrix is a function of the number of features included in the feature stream.

In some implementations, the first hidden layer 322 includes a number of LSTM logic units 322a. In some implementations, the number of LSTM logic units 322a ranges between approximately 10-500. Those of ordinary skill in the art will appreciate that, in such implementations, the number of LSTM logic units per layer is orders of magnitude smaller than previously known approaches (being of the order of $O(10^1)$ - $O(10^2)$), which allows such implementations to be embedded in highly resource-constrained devices. As illustrated in the example of FIG. 3B, the first hidden layer 322 receives its inputs from the input layer 320.

In some implementations, the second hidden layer 324 includes a number of LSTM logic units 324a. In some implementations, the number of LSTM logic units 324a is the same as or similar to the number of LSTM logic units 320a in the input layer 320 or the number of LSTM logic units 322a in the first hidden layer 322. As illustrated in the example of FIG. 3B, the second hidden layer 324 receives its inputs from the first hidden layer 322. Additionally or alternatively, in some implementations, the second hidden layer 324 receives its inputs from the input layer 320.

In some implementations, the classification layer **326** includes a number of LSTM logic units **326a**. In some implementations, the number of LSTM logic units **326a** is the same as or similar to the number of LSTM logic units **320a** in the input layer **320**, the number of LSTM logic units **322a** in the first hidden layer **322**, or the number of LSTM logic units **324a** in the second hidden layer **324**. In some implementations, the classification layer **326** includes an implementation of a multinomial logistic function (e.g., a soft-max function) that produces a number of candidate plans **314a** . . . **314n**. In some implementations, the number of candidate plans **314a** . . . **314n** is approximately equal to the number of possible plans **360**. In some implementations, the candidate plans **314a** . . . **314n** are associated with corresponding confidence scores **316a** . . . **316n** which include a probability or a confidence measure that the corresponding plan satisfies the objective **254**. In some implementations, the outputs do not include plans that have been excluded by operation of the limiter **370**.

In some implementations, the classification layer **326** includes the future outcome determiner **348**. In such implementations, the future outcome determiner **348** determines potential future outcomes for each of the candidate plans **314a** . . . **314n**. In some implementations, the classification layer **326** determines the confidence scores **316a** . . . **316n** for the candidate plans **314a** . . . **314n** based on the respective potential future outcomes of the candidate plans **314a** . . . **314n**. In some implementations, the classification layer **326** assigns a higher confidence score to candidate plans that result in desirable potential future outcomes, and a lower confidence score to candidate plans that result in undesirable potential future outcomes. In some implementations, a desirable potential future outcome refers to an outcome in which the objective for the objective-effectuator is satisfied, whereas an undesirable potential future outcome refers to an outcome in which the objective for the objective-effectuator is not satisfied.

In some implementations, the plan selector **328** generates the plan **314** by selecting the top N candidate plans provided by the classification layer **326**. For example, in some implementations, the plan selector **328** selects the candidate plan with the highest confidence score. In some implementations, the top N candidate plans are most likely to satisfy the objectives **254**. In some implementations, the plan selector **328** provides the plan **314** to a rendering and display pipeline (e.g., the display engine **260** shown in FIG. 2).

FIG. 3C illustrates a behavior tree (BT) **380**. In some implementations, the BT **380** includes a probabilistic behavior tree (PBT). In the example of FIG. 3C, the BT **380** generates the plan **314** based on one or more inputs. As illustrated in FIG. 3C, in some implementations, the inputs to the BT **380** include indications of the set of possible actions **209**, the instantiated equipment **340**, the instantiated characters **342**, the user-specified scene/environment information **344**, the objectives **254**, the characteristic value(s) **345**, the current state **346** of the objective-effectuator, the portion of the knowledge graph **347**, and/or the potential future outcomes **349**. In some implementations, searching/traversing the BT **380** with one or more of the inputs shown in FIG. 3C results in the plan **314**.

FIG. 4A is a flowchart representation of a method **400** of generating a plan for an objective-effectuator. In various implementations, the method **400** is performed by a device with a non-transitory memory and one or more processors coupled with the non-transitory memory (e.g., the controller **102** and/or the electronic device **103** shown in FIG. 1). In some implementations, the method **400** is performed by

processing logic, including hardware, firmware, software, or a combination thereof. In some implementations, the method **400** is performed by a processor executing code stored in a non-transitory computer-readable medium (e.g., a memory).

As represented by block **410**, in various implementations, the method **400** includes obtaining an objective for a CGR representation of an objective-effectuator. In some implementations, the objective is associated with a plurality of time frames (e.g., several seconds, minutes, hours, days, weeks, months, years or decades). In some implementations, the method **400** includes receiving the objective from an emergent content engine that generated the objective. For example, as shown in FIG. 2, the planner **204a** receives the objective **254** for the boy objective-effectuator, represented by the boy action figure representation **108a**, from the emergent content engine **250**. In some implementations, the method **400** includes generating the objective (e.g., by selecting the objective from a set of possible objectives). In some implementations, the method **400** includes retrieving the objective from a datastore.

As represented by block **420**, in various implementations, the method **400** includes determining a plurality of candidate plans that satisfy the objective. In some implementations, each of the plurality of candidate plans is associated with a corresponding confidence score. For example, as shown in FIG. 3B, the neural network **310** determines the candidate plans **314a** . . . **314n**, and each of the candidate plans **314a** . . . **314n** is associated with corresponding confidence scores **316a** . . . **316n**.

As represented by block **430**, in various implementations, the method **400** includes selecting a first candidate plan of the plurality of candidate plans based on a selection criterion (e.g., based on the respective confidence scores). For example, as shown in FIG. 3B, the plan selector **328** selects the plan **314** from the candidate plans **314a** . . . **314n** based on the confidence scores **316a** . . . **316n**. In some implementations, the method **400** includes selecting a candidate plan with a confidence score that satisfies a confidence threshold. For example, selecting a candidate plan with a confidence score that is greater than the confidence threshold. In some implementations, the method **400** includes selecting the candidate plan with the highest confidence score.

As represented by block **440**, in various implementations, the method **400** includes effectuating the first candidate plan in order to satisfy the objective. In some implementations, the first candidate plan triggers the CGR representation of the objective-effectuator to perform a series of actions over the plurality of time frames associated with the objective. In some implementations, the method **400** includes generation actions in accordance with the plan, and performing the generated actions. For example, as shown in FIG. 2, the planners **204a** . . . **204n** provide the plans **206a** . . . **206n** to the objective-effectuator engines in order to allow the objective-effectuator engines to generate the actions **210a** . . . **210n** in accordance with the plans **206a** . . . **206n**.

Referring to FIG. 4B, as represented by block **420a**, in some implementations, the method **400** includes determining the plurality of candidate plans based on one or more characteristic values associated with the objective-effectuator (e.g., the one or more characteristic values **345** shown in FIGS. 3A-3C). In some implementations, the characteristic values indicate various characteristics of the CGR representation of the objective-effectuator. In such implementations, the method **400** includes determining the plurality of candidate plans based on the characteristics of the CGR representation of the objective-effectuator. For example, the

method **400** includes assigning higher confidence scores to candidate plans that the characteristics support and assigning lower confidence scores to candidate plans that the characteristics do not support. In some implementations, the one or more characteristic values indicate one or more physical/structural characteristics of the CGR representation of the objective-effectuator (e.g., body material, accessories such as a jet pack, etc.). In some implementations, the one or more characteristic values indicate one or more behavioral characteristics of the CGR representation of the objective-effectuator (e.g., level of aggressiveness, mood, etc.). In some implementations, the one or more characteristic values indicate one or more functional characteristics of the CGR representation of the objective-effectuator (e.g., running, flying, defying gravity, etc.).

As represented by block **420b**, in some implementations, the method **400** includes utilizing a behavior tree (BT) (e.g., a probabilistic behavior tree (PBT) to generate the plurality of candidate plans (e.g., the BT **380** shown in FIG. 3C). In some implementations, the method **400** includes searching the BT for the plurality of candidate plans based on one or more characteristic values associated with the objective-effectuator. In some implementations, the method **400** includes traversing the BT for the plurality of candidate plans based on one or more characteristic values associated with the objective-effectuator. For example, in some implementations, the BT generates the candidate plans **314a** . . . **314n** shown in FIG. 3B.

As represented by block **420c**, in some implementations, the method **400** includes utilizing a reactive neural network to generate the plurality of candidate plans. For example, in some implementations, the reactive neural network generates the candidate plans **314a** . . . **314n** shown in FIG. 3B.

As represented by block **420d**, in some implementations, the method **400** includes generating the candidate plans based on a current state of the CGR representation of the objective-effectuator (e.g., based on the current state **346** shown in FIGS. 3A-3C). In some implementations, the method **400** includes obtaining a current state of the CGR representation of the objective-effectuator, and determining the plurality of candidate plans based on the current state of the CGR representation of the objective-effectuator. In some implementations, the method **400** includes utilizing an emphatic engine to evaluate the current state of the CGR representation of the objective-effectuator. In some implementations, the method **400** includes evaluating the current state of the CGR representation of the objective-effectuator based on a body pose of the CGR representation and/or based on actions being performed by the CGR representation. In some implementations, the method **400** includes providing the current state of the CGR representation of the objective-effectuator as an input to a neural network that generates the plurality of candidate plans (e.g., providing the current state as an input to the neural network **310** shown in FIG. 3A).

As represented by block **420e**, in some implementations, the objective-effectuator determines the plurality of candidate plans. In some implementations, the objective-effectuator includes a planner (e.g., the planner **300** shown in FIG. 3A) that determines the plurality of candidate plans. In some implementations, the planner is integrated into an objective-effectuator engine that generates actions for the CGR representation of the objective-effectuator (e.g., in some implementations, the planner **204a** is integrated into the character engine **208a** shown in FIG. 2).

As represented by block **420f**, in some implementations, the method **400** includes determining the candidate plans

based on known knowledge (e.g., based on the portion of the knowledge graph **347** shown in FIGS. 3A-3C). In some implementations, each objective-effectuator has access to a limited portion of a knowledge graph, and all objective-effectuators collectively have access to the entire knowledge graph. In such implementations, the method **400** includes determining the plurality of candidate plans based on the portion of the knowledge graph for which the objective-effectuator has access.

As represented by block **420g**, in some implementations, the method **400** includes generating the plurality of candidate plans based on reinforcement learning. For example, in some implementations, the method **400** includes assigning higher scores to desirable plans during a training phase, and assigning lower scores to undesirable plans during the training phase. In some implementations, the method **400** includes generating the plurality of candidate plans based on initiation learning.

Referring to FIG. 4C, as represented by block **430a**, in some implementations, the method **400** includes determining respective potential future outcomes for each of the plurality of candidate plans (e.g., the potential future outcomes **349** shown in FIGS. 3A-3C). For example, determining respective potential future outcomes for each of the candidate plans **314a** . . . **314n** shown in FIG. 3B.

As represented by block **430b**, in some implementations, the method **400** includes determining corresponding confidence scores for the plurality of candidate plans based on the potential future outcomes. For example, determining the confidence scores **316a** . . . **316n** for the candidate plans **314a** . . . **314n** shown in FIG. 3B. In some implementations, the method **400** includes assigning higher confidence scores to candidate plans that result in desirable potential future outcomes, and assigning lower confidence scores to candidate plans that result in undesirable potential future outcomes.

As represented by block **430c**, in some implementations, the method **400** includes ranking the plurality of candidate plans based on the respective potential future outcomes. For example, the method **400** includes assigning higher ranks to candidate plans that result in desirable potential future outcomes, and assigning lower ranks to candidate plans that result in undesirable potential future outcomes. In some implementations, the method **400** includes ranking the plurality of candidate plans based on the corresponding confidence scores.

As represented by block **430d**, in some implementations, the first candidate plan limits the series of actions to a permissible set of actions. In some implementations, the first candidate plan limits the series of actions in order to avoid an impermissible set of actions. In some implementations, the first candidate plan guides the objective-effectuator engine in generating the actions. For example, the first candidate plan indicates types of actions for the objective-effectuator engine to generate. In some implementations, the first candidate plan indicates a sequence for different types of actions for the objective-effectuator.

As represented by block **430e**, in some implementations, the method **400** includes modifying the BT (e.g., the PBT) based on the selection of the first candidate plan (e.g., modifying the BT **380** based on the selection of the plan **314**). In some implementations, the method **400** includes, after selecting the first candidate plan, assigning more weight (e.g., a higher probability) to a portion of the BT that corresponds to the first candidate plan. In some implementations, the method **400** includes, after selecting the first

candidate plan, assigning less weight (e.g., a lower probability) to a portion of the BT that does not correspond to the first candidate plan.

As represented by block **440a**, in some implementations, the method **400** includes displaying the CGR representation of the objective-effectuator performing the series of actions over the plurality of time frames in accordance with the first candidate plan. For example, displaying the boy action figure representation **108a** performing the actions **210a** that were generated in accordance with the plan **206a** shown in FIG. 2. In some implementations, the method **400** includes animating (e.g., manipulating) the CGR representation of the objective-effectuator in order to provide an appearance that the CGR representation of the objective-effectuator is performing the series of actions.

FIG. 5 is a block diagram of a device **500** enabled with one or more components of a planner (e.g., one of the planners **204a** . . . **204e** shown in FIG. 2, or the planner **300** shown in FIG. 3A) in accordance with some implementations. While certain specific features are illustrated, those of ordinary skill in the art will appreciate from the present disclosure that various other features have not been illustrated for the sake of brevity, and so as not to obscure more pertinent aspects of the implementations disclosed herein. To that end, as a non-limiting example, in some implementations the device **500** includes one or more processing units (CPUs) **501**, a network interface **502**, a programming interface **503**, a memory **504**, and one or more communication buses **505** for interconnecting these and various other components.

In some implementations, the network interface **502** is provided to, among other uses, establish and maintain a metadata tunnel between a cloud hosted network management system and at least one private network including one or more compliant devices. In some implementations, the one or more communication buses **505** include circuitry that interconnects and controls communications between system components. The memory **504** includes high-speed random access memory, such as DRAM, SRAM, DDR RAM or other random access solid state memory devices, and may include non-volatile memory, such as one or more magnetic disk storage devices, optical disk storage devices, flash memory devices, or other non-volatile solid state storage devices. The memory **504** optionally includes one or more storage devices remotely located from the one or more CPUs **501**. The memory **504** comprises a non-transitory computer readable storage medium.

In some implementations, the memory **504** or the non-transitory computer readable storage medium of the memory **504** stores the following programs, modules and data structures, or a subset thereof including an optional operating system **506**, a data obtainer **510**, a candidate plan determiner **520**, a plan selector **530**, and a plan effectuator **540**. In various implementations, the device **500** performs the method **400** shown in FIGS. 4A-4C.

In some implementations, the data obtainer **510** obtains an objective for a CGR representation of an objective-effectuator. In some implementations, the data obtainer **510** performs the operation(s) represented by block **410** in FIG. 4A. To that end, the data obtainer **510** includes instructions **510a**, and heuristics and metadata **510b**.

In some implementations, the candidate plan determiner **520** determines a plurality of candidate plans that satisfy the objective. In some implementations, the candidate plan determiner **520** performs the operation(s) represented by blocks **420** and **420a** . . . **420g** shown in FIGS. 4A-4B. To

that end, the candidate plan determiner **520** includes instructions **520a**, and heuristics and metadata **520b**.

In some implementations, the plan selector **530** selects a first candidate plan of the plurality of candidate plans based on the respective confidence scores. In some implementations, the plan selector **530** performs the operations represented by blocks **430** and **430a** . . . **430e** shown in FIGS. 4A and 4C. To that end, the plan selector **530** includes instructions **530a**, and heuristics and metadata **530b**.

In some implementations, the plan effectuator **540** effectuates the first candidate plan in order to satisfy the objective. In some implementations, the plan effectuator **540** performs the operations represented by blocks **440** and **440a** shown in FIGS. 4A and 4C. To that end, the plan effectuator **540** includes instructions **540a**, and heuristics and metadata **540b**.

While various aspects of implementations within the scope of the appended claims are described above, it should be apparent that the various features of implementations described above may be embodied in a wide variety of forms and that any specific structure and/or function described above is merely illustrative. Based on the present disclosure one skilled in the art should appreciate that an aspect described herein may be implemented independently of any other aspects and that two or more of these aspects may be combined in various ways. For example, an apparatus may be implemented and/or a method may be practiced using any number of the aspects set forth herein. In addition, such an apparatus may be implemented and/or such a method may be practiced using other structure and/or functionality in addition to or other than one or more of the aspects set forth herein.

It will also be understood that, although the terms “first”, “second”, etc. may be used herein to describe various elements, these elements should not be limited by these terms. These terms are only used to distinguish one element from another. For example, a first node could be termed a second node, and, similarly, a second node could be termed a first node, which changing the meaning of the description, so long as all occurrences of the “first node” are renamed consistently and all occurrences of the “second node” are renamed consistently. The first node and the second node are both nodes, but they are not the same node.

The terminology used herein is for the purpose of describing particular implementations only and is not intended to be limiting of the claims. As used in the description of the implementations and the appended claims, the singular forms “a”, “an”, and “the” are intended to include the plural forms as well, unless the context clearly indicates otherwise. It will also be understood that the term “and/or” as used herein refers to and encompasses any and all possible combinations of one or more of the associated listed items. It will be further understood that the terms “comprises” and/or “comprising”, when used in this specification, specify the presence of stated features, integers, steps, operations, elements, and/or components, but do not preclude the presence or addition of one or more other features, integers, steps, operations, elements, components, and/or groups thereof.

As used herein, the term “if” may be construed to mean “when” or “upon” or “in response to determining” or “in accordance with a determination” or “in response to detecting”, that a stated condition precedent is true, depending on the context. Similarly, the phrase “if it is determined [that a stated condition precedent is true]” or “if [a stated condition precedent is true]” or “when [a stated condition precedent is true]” may be construed to mean “upon determining” or “in

response to determining” or “in accordance with a determination” or “upon detecting” or “in response to detecting” that the stated condition precedent is true, depending on the context.

What is claimed is:

1. A method comprising:
at a device including a non-transitory memory, a display,
and one or more processors coupled with the non-transitory memory and the display:
displaying, on the display, a computer-generated reality (CGR) environment that is associated with state information that indicates one or more states of the CGR environment, wherein displaying the CGR environment is based on a first portion of the state information that is accessible to an objective-effectuator and a second portion of the state information that is not accessible to the objective-effectuator;
determining, based on the first portion of the state information that is accessible to the objective-effectuator and not the second portion of the state information that is not accessible to the objective-effectuator, a candidate plan that provides guidance for generating a series of actions that satisfies an objective of the objective-effectuator; and
while displaying the CGR environment, displaying, on the display, a CGR representation of the objective-effectuator performing the series of actions that is generated in accordance with the guidance provided by the candidate plan in order to advance the objective of the objective-effectuator.
2. The method of claim 1, wherein the one or more states include a current state of the CGR environment.
3. The method of claim 1, wherein the one or more states include a set of one or more historical states of the CGR environment.
4. The method of claim 1, wherein the CGR environment is associated with a time duration;
wherein the state information indicates an initial state of the CGR environment at a beginning of the time duration; and
wherein the candidate plan is a function of the initial state of the CGR environment at the beginning of the time duration.
5. The method of claim 4, wherein the initial state of the CGR environment indicates initial positions of CGR objects within the CGR environment at the beginning of the time duration; and
wherein the candidate plan is a function of the initial positions of the CGR objects at the beginning of the time duration.
6. The method of claim 4, wherein the initial state of the CGR environment indicates initial positions of virtual characters and virtual equipment in the CGR environment at the beginning of the time duration; and
wherein the candidate plan is a function of the initial positions of the virtual characters and the virtual equipment in the CGR environment.
7. The method of claim 1, wherein the CGR environment is associated with a time duration;
wherein the state information indicates an end state of the CGR environment at an end of the time duration; and
wherein the candidate plan is a function of the end state of the CGR environment at the end of the time duration.
8. The method of claim 7, wherein the end state of the CGR environment indicates end positions of CGR objects within the CGR environment at the end of the time duration; and

wherein the candidate plan is a function of the end positions of the CGR objects at the end of the time duration.

9. The method of claim 7, wherein the end state of the CGR environment indicates end positions of virtual characters and virtual equipment in the CGR environment at the end of the time duration; and

wherein the candidate plan is a function of the end positions of the virtual characters and the virtual equipment in the CGR environment.

10. The method of claim 1, wherein the state information indicates environmental conditions within the CGR environment; and

wherein determining the candidate plan comprises:

forgoing selection of a first candidate plan, from a plurality of candidate plans, that is adversely impacted by the environmental conditions within the CGR environment; and

selecting a second candidate plan, from the plurality of candidate plans, based on the environmental conditions within the CGR environment.

11. The method of claim 1, wherein determining the candidate plan comprises selecting the candidate plan from a plurality of candidate plans.

12. The method of claim 11, wherein selecting the candidate plan comprises selecting the candidate plan based on a selection criterion.

13. The method of claim 12, wherein selecting the candidate plan based on the selection criterion comprises selecting the candidate plan based on respective confidence scores of the plurality of candidate plans.

14. The method of claim 1, wherein displaying the CGR representation of the objective-effectuator corresponds to adding the CGR representation of the objective-effectuator to the display while displaying the CGR environment.

15. The method of claim 1, wherein the guidance indicates types of actions that can be generated for the objective-effectuator.

16. The method of claim 1, wherein the guidance indicates a permissible set of actions and an impermissible set of actions.

17. A device comprising:

one or more processors;
a non-transitory memory;
a display; and

one or more programs stored in the non-transitory memory, which, when executed by the one or more processors, cause the device to:

display, on the display, a computer-generated reality (CGR) environment that is associated with state information that indicates one or more states of the CGR environment, wherein displaying the CGR environment is based on a first portion of the state information that is accessible to an objective-effectuator and a second portion of the state information that is not accessible to the objective-effectuator;

determine, based on the first portion of the state information that is accessible to the objective-effectuator and not the second portion of the state information that is not accessible to the objective-effectuator, a candidate plan that provides guidance for generating a series of actions that satisfies an objective of the objective-effectuator; and

while displaying the CGR environment, display, on the display, a CGR representation of the objective-effectuator performing the series of actions that is gener-

25

ated in accordance with the guidance provided by the candidate plan in order to advance the objective of the objective-effectuator.

18. The device of claim 17, wherein the CGR environment is associated with a time duration;

wherein the state information indicates an initial state of the CGR environment at a beginning of the time duration; and

wherein the candidate plan is a function of the initial state of the CGR environment at the beginning of the time duration.

19. The device of claim 18, wherein the initial state of the CGR environment indicates initial positions of CGR objects within the CGR environment at the beginning of the time duration; and

wherein the candidate plan is a function of the initial positions of the CGR objects at the beginning of the time duration.

20. The device of claim 18, wherein the initial state of the CGR environment indicates initial positions of virtual characters and virtual equipment in the CGR environment at the beginning of the time duration; and

wherein the candidate plan is a function of the initial positions of the virtual characters and the virtual equipment in the CGR environment.

21. The device of claim 17, wherein the state information indicates environmental conditions within the CGR environment; and

wherein determining the candidate plan comprises:
 foregoing selection of a first candidate plan, from a plurality of candidate plans, that is adversely impacted by the environmental conditions within the CGR environment; and
 selecting a second candidate plan, from the plurality of candidate plans, based on the environmental conditions within the CGR environment.

26

22. A non-transitory memory storing one or more programs, which, when executed by one or more processors of a device with a display, cause the device to:

display, on the display, a computer-generated reality (CGR) environment that is associated with state information that indicates one or more states of the CGR environment, wherein displaying the CGR environment is based on a first portion of the state information that is accessible to an objective-effectuator and a second portion of the state information that is not accessible to the objective-effectuator;

determine, based on the first portion of the state information that is accessible to the objective-effectuator and not the second portion of the state information that is not accessible to the objective-effectuator, a candidate plan that provides guidance for generating a series of actions that satisfies an objective of the objective-effectuator; and

while displaying the CGR environment, display, on the display, a CGR representation of the objective-effectuator performing the series of actions that is generated in accordance with the guidance provided by the candidate plan in order to advance the objective of the objective-effectuator.

23. The non-transitory memory of claim 22, wherein the CGR environment is associated with a time duration;

wherein the state information indicates an initial state of the CGR environment at a beginning of the time duration; and

wherein the candidate plan is a function of the initial state of the CGR environment at the beginning of the time duration.

* * * * *