US Patent & Trademark Office
Patent Public Search | Text View

United States Patent	12395311
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Kumar; Seema et al.

Reliable link management for a high-speed signaling interconnect

Abstract

A device includes receiver circuitry to receive incoming signals on a clock lane and data lanes and detection circuitry. The detection circuitry is to monitor the incoming signals on the clock lane, and determine that an incoming pattern of the incoming signals on the clock lane does not correspond to a clock pattern associated with communicating data on the data lanes. The detection circuitry is to initiate a power-down sequence in response to determining that the incoming pattern does not correspond to the clock pattern.

Inventors: **Kumar; Seema (Santa Clara, CA), Chadha; Ish (San Jose, CA)**

Applicant: **NVIDIA Corporation** (Santa Clara, CA)

Family ID: **1000008765032**

Assignee: **NVIDIA Corporation (Santa Clara, CA)**

Appl. No.: **18/587111**

Filed: **February 26, 2024**

Prior Publication Data

Document Identifier	Publication Date
US 20240195599 A1	Jun. 13, 2024

Related U.S. Application Data

continuation parent-doc US 17988551 20221116 US 11956342 child-doc US 18587111
us-provisional-application US 63294029 20211227

Publication Classification

Int. Cl.: **H04L7/00** (20060101)

U.S. Cl.:

CPC   **H04L7/0083** (20130101); **H04L7/0012** (20130101);

Field of Classification Search

CPC: H04L (7/0083); H04L (7/0012); H04L (25/14); H04L (7/0008); H04L (43/08); H04L (12/12); G06F (1/26); G06F (1/266); G06F (1/3209)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
9496879	12/2015	Duan et al.	N/A	N/A
2008/0204561	12/2007	Yang	348/194	G06F 1/3203
2016/0147684	12/2015	Sengoku	N/A	N/A

OTHER PUBLICATIONS

Turner et al. "Ground-referenced signaling for intra-chip and short-reach chip-to-chip interconnects." 2018 IEEE Custom Integrated Circuits Conference (CICC)—Apr. 2018 (Year: 2018). cited by examiner

---

Primary Examiner: Huang; David S

Attorney, Agent or Firm: Lowenstein Sandler LLP

Background/Summary

RELATED APPLICATION (1) This application is a continuation of Ser. No. 17/988,551, filed on Nov. 16, 2022, which claims the benefit of U.S. Provisional Patent Application No. 63/294,029, filed Dec. 27, 2021, which are incorporated by reference herein in their entirety.

TECHNICAL FIELD
(1) At least one embodiment pertains to processing resources used to perform and facilitate high-speed communications. For example, at least one embodiment pertains to technology for reliable link management for a ground-referenced signaling (GRS) interconnect.
BACKGROUND
(2) Communication systems transmit signals from a transmitter to a receiver via a communication channel or medium (e.g., cables, printed circuit boards, links, wirelessly, etc.) For example, the communication channel can communicate signals between devices or chips—e.g., a chip to chip (C2C) system. In some communication systems, the communication channel at one of the devices or chips can be powered down or reset by a respective software stack for fault containment or a controlled shutdown operation. Each chip can utilize a different software stack, so the respective software stack of the powered down or reset device cannot communicate with the other chip. Additionally, the software stack can be unable to control the physical communication channel itself due to security specifications of the communication system. This can cause the device to shut down or reset abruptly. When the communication channel shuts down abruptly, it can fail to follow

proper power-down sequences causing additional stress on the communication channel and reducing the reliability of analog circuits of the communication channel over time.

Description

BRIEF DESCRIPTION OF DRAWINGS

(1) Various embodiments in accordance with the present disclosure will be described with reference to the drawings, in which:

(2) FIG. **1** is an example communication system employing a method for reliable link management, in accordance with at least some embodiments;

(3) FIG. **2** illustrates an example communication system employing a method for reliable link management for a high-speed interconnect, in accordance with at least some embodiments;

(4) FIG. **3** is an example communication system employing a method for reliable link management, in accordance with at least some embodiments;

(5) FIG. **4** is an example receiver employing a method for reliable link management for a high-speed interconnect system, in accordance with at least some embodiments;

(6) FIG. **5** is an example communication system employing a method for reliable link management, in accordance with at least some embodiments;

(7) FIG. **6** is a flow diagram of a method for reliable link management for a high-speed interconnect system, in accordance with at least some embodiments;

(8) FIG. **7** illustrates an example computer system including a transceiver including a chip-to-chip interconnect for reliable link management, in accordance with at least some embodiments.

DETAILED DESCRIPTION

(9) Communication systems transmit signals from a transmitter to a receiver via a communication channel or medium (e.g., cables, printed circuit boards, links, wirelessly, etc.). Some communication systems can include multiple devices executing separate, isolated software stacks. For example, the communication system can include a first device (e.g., a first integrated circuit (IC) or chip) and a second device (e.g., a second IC or chip) and communicate data via a ground-referenced signaling (GRS) link—e.g., the communication system may be a chip-to-chip (C2C) interconnect with both devices including a transmitter and a receiver. The first device and second device can execute isolated software stacks which can cause the first device and second device to be brought out of reset asynchronously—e.g., the first device or second device can be shut down or powered on at any time without the other device knowing. In some communication systems, the link can be powered-down or reset at either the first device or the second device by their respective software stack. Because the software stacks are isolated, the software stack powering down or resetting the link cannot communicate the power down or reset to the other device. Additionally, the respective software stack cannot access local physical link controls due to system security specifications. Accordingly, the power down and reset can be abrupt.

(10) Some communication systems can shut down the link and components associated with the link (e.g., transmitter and receiver components) according to a power-down sequence. That is, components associated with and of the link can have a voltage range (e.g., below a maximum voltage threshold) within which they operate reliably based on a manufacturing process. Exceeding the voltage range can cause additional stress on the components and cause them to break down over time. For example, the communication system can shut down certain analog components of the transmitter and receiver according to the power-down sequence. When components are shut down or reset out of sequence (e.g., not according to the power-down sequence), long-term reliability of the components can be reduced. For example, if the power-down sequence is not followed and the receiver termination is disabled at the first device while a transmit equalization is enabled at the second device, the transmitter can undergo additional stress, and the performance of the transmitter

is degraded. Additionally, if the first device stops transmitting a clock signal, a phase-lock loop (PLL) of the second device can begin to drift (e.g., become out of phase) as the PLL relies on the clock signal as a reference—e.g., relies on a forwarded clock signal for reference. When the shutdown or reset of the components is abrupt, the components associated with and of the link can be shutdown or reset out of sequence—e.g., when the respective software stack is unable to communicate the shutdown or reset to its local link components or the other device, the shutdown or reset can be out of sequence. Accordingly, the voltage range of some components can be exceeded and cause reliability issues.

(11) Advantageously, aspects of the present disclosure can address the deficiencies above and other challenges by providing a method for reliable link management by detecting a status of a link coupling a first device and a second device via a clock lane. For example, each device can utilize link status detection circuitry in their receivers to monitor incoming signals and patterns on the clock lane. In some embodiments, each device can configure counter logic to determine a number of pulses received over a configurable period of time—e.g., the device can determine the average frequency of incoming patterns by detecting edges in the pattern for the period of time. If the device determines the incoming pattern does not correspond to a clock pattern associated with transmitting data (e.g., the number of pulses determined differs from an expected number of pulses associated with a clock signal for transmitting data), the device can determine the link or its associated components is down at the other side—e.g., at the other device. For example, a software stack of the first device can initiate a shutdown or reset of the link at the first device. In such embodiments, the first device can either begin to transmit a static pattern on the clock lane to indicate the first device is shutting down or resetting the link, or the first device can shutdown or reset the transmitter and stop driving a signal on the clock lane. In either case, the second device can continuously determine a number of pulses received over a configurable period of time. When the first device transmits the static pattern or stops driving a signal on the clock lane, the second device can determine if the number of pulses received over the configurable period does not match the expected number of pulses associated with a clock signal for transmitting data. Accordingly, the second device can determine the link is being powered down or reset at the first device and initiate its own power-down sequence. In some embodiments, the device can set longer or shorter periods of time based on how quickly the power-down sequence is to be initiated—e.g., if a transmitter of the second device has to be shutdown within a first period of time after a receiver of the first device is shutdown, the second device can set the period of time to be less than the first period.

(12) Additionally, each device can include ramp-down (e.g., power down) circuitry configured to receive a link status—e.g., an indication of whether the determined number of pulses satisfies the expected number of pulses. The ramp-down circuitry can enable the device to initiate the power-down sequence when the device determines the link is powering down or resetting at the other device—e.g., the ramp-down circuitry can control the physical link controls to proceed with the power-down sequence. For example, when the device determines the link is powering down or resetting at the other device, the ramp-down circuitry can receive the indication and proceed to disable the transmitter equalization and the PLL of the receiver to maintain the reliability of the components. The ramp-down circuitry can also be configured to receive a reset signal from error circuitry or an indication from the respective software stack to power down and reset. Accordingly, even if the software stack does not have control of the physical link controls, the software stack can indicate the power down to the ramp-down circuitry, which can access the controls of the link and initiate the power-down sequence.

(13) By utilizing the clock lane to monitor the link status, each device can properly execute the power-down sequence. That is, a device can quickly initiate a power-down sequence after determining the link is powering down or resetting at the other device without software intervention. Additionally, the solution can avoid circuits design with no reliability loopholes—e.g., circuits with no reliability loopholes can be costly to design and still cannot account for all

situations. Accordingly, embodiments of the present application allow for a hardware-autonomous method for reliable link management (e.g., more reliable power down or resets of the link) in a high-speed interconnect system without the added design costs.

(14) FIG. **1** illustrates an example communication system **100** according to at least one example embodiment. The system **100** includes a device **110**, a communication network **108** including a communication channel **109**, and a device **112**. In at least one embodiment, devices **110** and **112** are two end-point devices in a computing system, such as a central processing unit (CPU) or graphics processing unit (GPU). In at least one embodiment, devices **110** and **112** are two servers. In at least one example embodiment, devices **110** and **112** correspond to one or more of a Personal Computer (PC), a laptop, a tablet, a smartphone, a server, a collection of servers, or the like. In some embodiments, the devices **110** and **112** may correspond to any appropriate type of device that communicates with other devices connected to a common type of communication network **108**. According to embodiments, the receiver **104** of devices **110** or **112** may correspond to a GPU, a switch (e.g., a high-speed network switch), a network adapter, a CPU, a memory device, an input/output (I/O) device, other peripheral devices or components on a system-on-chip (SoC), or other devices and components at which a signal is received or measured, etc. As another specific but non-limiting example, the devices **110** and **112** may correspond to servers offering information resources, services, and/or applications to user devices, client devices, or other hosts in the system **100**. In one example, devices **110** and **112** may correspond to network devices such as switches, network adapters, or data processing units (DPUs).

(15) Examples of the communication network **108** that may be used to connect the devices **110** and **112** include an Internet Protocol (IP) network, an Ethernet network, an InfiniBand (IB) network, a Fibre Channel network, the Internet, a cellular communication network, a wireless communication network, a ground referenced signaling (GRS) link, combinations thereof (e.g., Fibre Channel over Ethernet), variants thereof, and/or the like. In one specific but non-limiting example, the communication network **108** is a network that enables data transmission between the devices **110** and **112** using data signals (e.g., digital, optical, wireless signals).

(16) The device **110** includes a transceiver **116** for sending and receiving signals, for example, data signals. The data signals may be digital or optical signals modulated with data or other suitable signals for carrying data.

(17) The transceiver **116** may include a digital data source **120**, a transmitter **124**, a receiver **104**, and processing circuitry **132** that controls the transceiver **116**. The digital data source **120** may include suitable hardware and/or software for outputting data in a digital format (e.g., in binary code and/or thermometer code). The digital data output by the digital data source **120** may be retrieved from memory (not illustrated) or generated according to input (e.g., user input).

(18) The transmitter **124** includes suitable software and/or hardware for receiving digital data from the digital data source **120** and outputting data signals according to the digital data for transmission over the communication network **108** to a receiver **104** of device **112**. Additional details of the structure of the transmitter **124** are discussed in more detail below with reference to the figures.

(19) The receiver **104** of devices **110** and **112** may include suitable hardware and/or software for receiving signals, such as data signals from the communication network **108**. For example, the receiver **104** may include components for receiving processing signals to extract the data for storing in a memory, as described in detail below with respect to FIG. **2**-FIG. **5**. In at least one embodiment, receiver **104** can include link status detection circuitry **115**. In some embodiments, link status detection circuitry **115** can be configured to receive clock signals from clock lanes of the communication network **108**. In some embodiments, link status detection circuitry **115** can determine the status of the communication network **108**—e.g., the link status detection circuitry **115** can determine whether the link is powered on at the other device. For example, the link status detection circuitry **115** is configured to receive a clock pattern and determine a number of pulses detected over a pre-defined period. If the link status detection circuitry **115** determines the number

of pulses detected satisfies an expected number of pulses, the device **110** or device **112** can determine the link is powered on at the other device. If the link status detection circuitry **115** determines the number of pulses detected fails to satisfy the expected number of pulses, the device **110** or device **112** can determine the link is powered off or reset at the other device. In such embodiments, the link status detection circuitry **115** can generate an indication (e.g., a link status) that the link is powered down at the other device and transmit the indication to a ramp-down component as described with reference to FIG. **5**. The ramp-down component can initiate a power-down sequence in response to receiving the indication, enabling the device **110** or device **112** to follow the power-down sequence and maintain the reliability of the link components and components associated with the link.

(20) The processing circuitry **132** may comprise software, hardware, or a combination thereof. For example, the processing circuitry **132** may include a memory including executable instructions and a processor (e.g., a microprocessor) that executes the instructions on the memory. The memory may correspond to any suitable type of memory device or collection of memory devices configured to store instructions. Non-limiting examples of suitable memory devices that may be used include Flash memory, Random Access Memory (RAM), Read Only Memory (ROM), variants thereof, combinations thereof, or the like. In some embodiments, the memory and processor may be integrated into a common device (e.g., a microprocessor may include integrated memory). Additionally or alternatively, the processing circuitry **132** may comprise hardware, such as an application-specific integrated circuit (ASIC). Other non-limiting examples of the processing circuitry **132** include an Integrated Circuit (IC) chip, a Central Processing Unit (CPU), a General Processing Unit (GPU), a microprocessor, a Field Programmable Gate Array (FPGA), a collection of logic gates or transistors, resistors, capacitors, inductors, diodes, or the like. Some or all of the processing circuitry **132** may be provided on a Printed Circuit Board (PCB) or collection of PCBs. It should be appreciated that any appropriate type of electrical component or collection of electrical components may be suitable for inclusion in the processing circuitry **132**. The processing circuitry **132** may send and/or receive signals to and/or from other elements of the transceiver **116** to control the overall operation of the transceiver **116**.

(21) The transceiver **116** or selected elements of the transceiver **116** may take the form of a pluggable card or controller for the device **110**. For example, the transceiver **116** or selected elements of the transceiver **116** may be implemented on a network interface card (NIC).

(22) The device **112** may include a transceiver **136** for sending and receiving signals, for example, data signals over a channel **109** of the communication network **108**. The same or similar structure of the transceiver **116** may be applied to transceiver **136**, and thus, the structure of transceiver **136** is not described separately.

(23) Although not explicitly shown, it should be appreciated that devices **110** and **112** and the transceivers **116** and **136** may include other processing devices, storage devices, and/or communication interfaces generally associated with computing tasks, such as sending and receiving data.

(24) FIG. **2** illustrates an example communication system **200** according to at least one example embodiment. The system **200** includes a device **110** and a device **112** as described with reference to FIG. **1**. Device **110** and device **112** can include a transmitter **124** and receiver **104** as described with reference to FIG. **1**. The device **110** and device **112** can be coupled to a link **220**. In at least one embodiment, the link **220** can be an example of communication network **108** as described with reference to FIG. **1**. In at least one embodiment, link **220** can be an example of a high-speed interconnect. For example, link **220** can be an example of a forwarded clock architecture. In one embodiment, the link **220** can be an example of a ground referenced signaling (GRS) link **220**. In an embodiment, the GRS link **220** can be a signaling scheme used for serial data transfer between devices **110** and **112**. In at least one embodiment, the GRS link **220** can be a high-speed link (e.g., transferring 40 gigabits per second (GBPS) at a frequency of 20 gigahertz when performing high-

speed communications). In at least one embodiment, the link **220** may include RC-dominated channels and LC transmission lines. Additionally, the GRS link **220** may be an on-chip link, a link across a substrate (e.g., organic package), or link signaling over a printed circuit board (PCB). In some examples, GRS link **220** may use a ground network as a signal reference voltage—e.g., the ground may be the return signaling.

(25) In at least one embodiment, the link **220** can include data lanes **202** and data lanes **203** configured to transmit signals, data, messages, etc., between the device **110** and device **112**. For example, data lanes **202** can be associated with communicating signals, data, or messages from device **110** to device **112**. Data lanes **203** can be associated with communicating signals, data, or messages from device **112** to device **110**—e.g., data lanes **202** can be associated with a transmitter **124** of device **110**, and data lanes **203** can be associated with a transmitter of device **112**. In at least one embodiment, the link **220** can include a same number of data lanes **202** and data lanes **203**. In this embodiment, a data lane **202** can be associated with a data lane **203**—e.g., data lane **202**-*a* and data lane **203**-*a* can be a single transmitter/receiver data lane pair. In at least one embodiment, the link **220** can include an "N" number of data lane pairs—e.g., an "N" number of data lanes **202** and data lanes **203**. In some embodiments, data lanes **202** can be associated with a forwarded clock lane **205**-*a*, and data lanes **203** can be associated with a forwarded clock lane **205**-*b*. In at least one embodiment, each clock lane can be associated with two or more data lanes—e.g., at least two data lanes **202** or data lanes **203**. In at least one embodiment, data lanes **202** transmit data to device **112**. In such embodiments, the data is latched on the forwarded clock at the receiver **104** of device **112**. In some embodiments, data lane **202** and the corresponding data lane **203** are identical—e.g., each data lane **202** and the corresponding data lane **203** support the same signaling speed and include identical drivers and hardware.

(26) In at least one embodiment, device **110** and device **112** can include link status detection circuitry **115**. In some embodiments, the link status detection circuitry **115** can be coupled with a clock lane **205**—i.e., link status detection circuitry **115**-*a* can be coupled with clock lane **205**-*a*, and link status detection circuitry **115**-*b* can be coupled with clock lane **205**-*b*. In some embodiments, a respective software stack of device **110** or device **112** can initiate a shutdown or reset of the link **220** on its side—e.g., the software stack of device **110** can shut down or reset a receiver and transmitter of device **110**. In some embodiments, devices **110** and devices **112** can execute isolated software stacks—e.g., the software stack of device **110** can be unable to communicate with the software stack of device **112**. To ensure a proper shutdown or power down, device **110** or device **112** has to ascertain whether the link **220** at the other device is powering down or reset. That is, components of the receiver and transmitter of devices **110** and device **112** can follow a power down (e.g., ramp down) sequence to ensure the components do not exceed a respective voltage threshold. Accordingly, device **110** and device **112** can utilize the link status detection circuitry **115** to determine a status of the link **220**, and when the link **220** is powered down or reset at the other device, initiate the power-down sequence. For example, the link status detection circuitry **115**-*a* can receive a clock pattern or signal on the respective clock lane **205**-*a*. The link status detection circuitry **115** can determine an average frequency of the clock pattern or signal received. In some embodiments, if the link status detection circuitry **115** determines the average frequency of the incoming clock pattern or signal satisfies an expected average frequency of a clock signal associated with transmitting data over the link **220**, the link status detection circuitry **115** can determine the link **220** at the other device is active. In such embodiments, device **110** and device **112** can continue transmitting and receiving data. In some embodiments, if the link status detection circuitry **115** determines the average frequency of the incoming clock pattern or signal fails to satisfy an expected average frequency of a clock signal associated with transmitting data over the link **220**, the link status detection circuitry **115** can determine the link **220** at the other device is shut down, reset, being shut down, or being reset. In such embodiments, the link status detection circuitry **115** can transmit an indication to a ramp-down component as described with reference to

FIG. **5**. Accordingly, the respective device **110** or device **112** can initiate a power-down sequence and ensure the components are shut down reliably without exceeding their respective threshold voltages.

(27) FIG. **3** illustrates an example communication system **300** according to at least one example embodiment. In at least one embodiment, communication system **300** is an example of communication system **100** or **200** as described with reference to FIGS. **1** and **2**. The system **300** includes a receiver **104** as described with reference to FIG. **1**. The receiver **104** can be in device **110** or device **112**. In some embodiments, the receiver **104** can be coupled with a link **220** as described with reference to FIG. **2**. The receiver **104** can include link status detection circuitry **115** and reliability hardware **330**.

(28) In some embodiments, receiver **104** can be configured to determine a link status **325** of link **220**. For example, receiver **104** can monitor incoming clock signals **305** received on a clock lane (e.g., clock lane **205** as described with reference to FIG. **2**) to determine the link status **325** of link **220** at the other device—e.g., the receiver **104** of device **110** can determine the link status **325** at device **110** by monitoring the incoming clock signal **305**. In some embodiments, the link status **325** can indicate whether the link **220** is active or inactive at the other device—e.g., whether the link **220** is driving data or being powered down or reset.

(29) In some embodiments, the link status detection circuitry **115** can receive the clock pattern **310** from the clock lane **205**. In some embodiments, the clock pattern **310** is associated with transmitting data or an active link **220**. For example, during normal operation, a transmitter (e.g., transmitter **124** as described with reference to FIG. **2**) can transmit a periodic clock signal **305** on the clock lane **205**, where the periodic clock signal **305** is the clock pattern **310**. In some embodiments, the link status detection circuitry **115** can determine (e.g., detect) a number of pulses (e.g., a number of rising or falling edges) over a configurable period for an incoming clock signal **305**. In at least one embodiment, the link status detection circuitry **115** can determine a number of pulses detected during the period that satisfies an expected number of pulses (e.g., satisfies a predetermined condition relating to a specified number of pulses for the period) when receiving the clock pattern **310**. That is, link status detection circuitry **115** can be programmed to compare the detected number of pulses during the period with an expected number of pulses associated with an active link **220**—e.g., associated with the periodic clock pattern **310**. Because the clock pattern **310** is associated with normal operation (e.g., an active link), the link status detection circuitry **115** can determine the number of pulses detected while receiving the clock pattern **310** satisfies the expected number of pulses. In such embodiments, the link status detection circuitry **115** can transmit a link status **325**, indicating the link **220** is active to the reliability hardware **330**. In some embodiments, when the reliability hardware **330** receives the link status **325** indicating the link **220** is active, the reliability hardware **330** can refrain from taking an additional action **335**. In at least one embodiment, the reliability hardware **330** is an example of a finite state machine (FSM) configured to initiate and manage a power-down sequence as described with reference to FIG. **5**.

(30) In some embodiments, the link status detection circuitry **115** can receive a static clock **315** from the clock lane **205**. In some embodiments, the static clock **315** is associated with a device **110** or device **112**, indicating it is about to be reset or powered down. For example, a software stack of device **110** can initiate a reset or shutdown of device **110**. In such embodiments, device **110** can begin driving a static clock **315** on the clock lane **205** to indicate to device **112** that device **110** is being reset or shut down. In at least one embodiment, the link status detection circuitry **115** can determine a number of pulses detected during the period that fails to satisfy an expected number of pulses (e.g., fails to satisfy the predetermined condition relating to the specified number of pulses for the period) when receiving the static clock **315**. That is, the static clock **315** can have a different number of pulses than the expected number of pulses associated with the clock pattern **310**. In such embodiments, the link status detection circuitry **115** can transmit a link status **325**, indicating the link **220** is inactive to the reliability hardware **330**. In some embodiments, when the reliability

hardware **330** receives the link status **325** indicating the link **220** is inactive, the reliability hardware **330** can take an action **335** to initiate the power-down sequence as described with reference to FIG. **5**.

(31) In some embodiments, the link status detection circuitry **115** can receive a disabled clock **320** (e.g., an un-driven clock **320**) from the clock lane **205**. In some embodiments, the disabled clock **320** is associated with device **110** or device **112** not driving anything across the clock lane **205**. Because nothing is driven on the clock lane **205**, the link status detection circuitry can receive random noise or toggles. In some embodiments, when the respective software stack of the device initiates the power down or reset of the link **220** or the associated components, the device can be unable to transmit the static clock **315** to indicate the shutdown. For example, the device **110** can ramp down or power down its transmitter **124** before the transmitter **124** can transmit the static clock **315**. In such embodiments, there is no clock signal **305** driven on clock lane **205**-*a*. Accordingly, the receiver **104** of device **112** can receive the disabled clock **320**, including random noise/toggles. In at least one embodiment, the link status detection circuitry **115** can determine a number of pulses detected during the period that fails to satisfy an expected number of pulses (e.g., fails to satisfy a predetermined condition relating to a specified number of pulses for the period) when receiving the disabled clock **320**. That is, the disabled clock **320** can have a different number of pulses (e.g., zero (0)) than the expected number of pulses associated with the clock pattern **310**. In some embodiments, the link status detection circuitry **115** does not determine whether the static clock **315** is received or the disabled clock **320** is received—e.g., the link status detection circuitry **115** determines that the incoming signal is different than the clock pattern **310** rather than determining what specific pattern is received. In some embodiments, the link status detection circuitry **115** can transmit a link status **325**, indicating the link **220** is inactive to the reliability hardware **330** when receiving the disabled clock **320**. As described above, when the reliability hardware **330** receives the link status **325** indicating the link **220** is inactive, the reliability hardware **330** can take an action **335** to initiate the power-down sequence as described with reference to FIG. **5**.

(32) FIG. **4** illustrates an example communication system **400** for reliable link management according to at least one embodiment. In at least one embodiment, communication system **400** is an example of communication system **100** or **200** as described with reference to FIGS. **1** and **2**. The system **400** includes a device **112** as described with reference to FIG. **1**. The device **112** can be coupled to a link **220** as described with reference to FIG. **2**—e.g., a GRS link **220**. Device **112** can include a receiver **104** as described with reference to FIG. **1**. In at least one embodiment, the components illustrated in receiver **104** can be considered a part of link status detection circuitry **115** as described with reference to FIG. **1**. Although FIG. **4** illustrates circuitry corresponding to clock lane **205**-*a*, communication system **400** can include similar circuitry corresponding to clock lane **205**-*b* as described with reference to FIG. **2**. That is, device **110** can also include a receiver **104**. Receiver **104** can include sampler **405**, dividers **410**, multiplexer **415**, low pass filter (LPF) **425**, local phase lock loop (PLL) **420**, and clock counter **430**. In some embodiments, the LPF **425**, local PLL **420**, and clock counter **430** can be digital logic. In at least one embodiment, clock counter **430** can be coupled to a controller or otherwise controlled by signals received from processing circuitry **132** as described with reference to FIG. **1** or receive signals from a respective software stack of device **112**.

(33) As described with reference to FIGS. **2** and **3**, in some embodiments, a respective software stack of a device **110** coupled to device **112** can initiate a shutdown or reset of the link **220** at the device **110**. To avoid having components of device **112** exceed voltage thresholds (e.g., to ensure a power-down sequence of device **112** is executed), the receiver **104** is configured to determine a link status **325** as described with reference to FIG. **3**—e.g., determine whether the link **220** is active at the device **110**.

(34) In an embodiment, sampler **405** can be configured to receive a set of bits corresponding to a

clock signal (e.g., clock signal **305** as described with reference to FIG. **3**) and sample the set of bits. For example, the sampler **405** can receive a first set of bits corresponding to clock pattern **310**, a second set of bits corresponding to static clock **315**, or a third set of bits corresponding to disabled clock **320** as described with reference to FIG. **3**. In at least one embodiment, sampler **405** can transmit the sampled set of bits received to dividers **410**. In some embodiments, sampler **405** can transmit the sampled set of bits to the multiplexer **415**—e.g., refrain from transmitting the sampled set of bits to the dividers **410**. In some embodiments, divider **410** can be configured to divide down the frequency of the set of bits received from device **110**. For example, divider **410** can be configured to divide a clock speed received by 2, 4, 8, 16, etc. In some embodiments, device **112** can configure the dividers **410** during an initialization of the receiver **104**. In some embodiments, the device **112** can further divide the clock speed received from device **110** to ensure the clock signal **305** received from device **110** is reliable. Although three dividers **410** is shown, receiver **104** can include any number of dividers **410** (e.g., include 1, 2, 4, 8, 16, etc.) based on the preferred clock speed of the device **112**. In some embodiments, a number of dividers **410** used can be configured by a software stack of device **112**. Dividers **410** can be configured to transmit the divided clock signal **305** (e.g. the divided sampled set of bits) to the multiplexer **415**. In some embodiments, the multiplexer **415** is configured to multiplex the set of bits received from the sampler **405** or the divided set of bits received from dividers **410** and transmit the multiplexed bits to the low pass filter **425**.

(35) In an embodiment, low pass filter **425** can be configured to refine further the sampled set of bits (or divided samples) from the multiplexer **415**. That is, the low pass filter **425** can remove short-term fluctuations and reduce noise on the sampled set of bits to enable the clock counter **430** to receive an improved sample or signal. In at least one embodiment, the low pass filter **425** can be coupled to a local PLL **420**. In some embodiments, the local PLL **420** can generate output signals for the low pass filter **425**. For example, the local PLL **420** can be initialized or configured to operate at the high-speed frequency (e.g., the frequency used for data transmission operations). In such embodiments, the local PLL **420** can be used to sample the received samples at the low pass filter **425**. In that, the local PLL **420** and low pass filter **425** can be utilized to get rid of high-frequency components in an output of the received clock—e.g., get rid of the higher frequencies to produce a signal with reduced noise for the clock counter **430**.

(36) In an embodiment, clock counter **430** can be configured to detect a number of pulses in the sampled set of bits (e.g., the received clock signal **305**) received from the low pass filter **425** during a pre-defined period—e.g., or from the multiplexer **415** in embodiments where a low pass filter **425** is not used. In some embodiments, clock counter **430** can detect the number of pulses by determining a number of edges (e.g., the number of times the signal rises or falls) during the pre-defined duration. In that, the clock counter **430** can determine the average frequency of the sampled set of bits—e.g., determine the average frequency of the clock signal **305** received at the receiver **104** from the device **110**. In some embodiments, the pre-defined period can be configurable. For example, the processing circuitry **132**, a controller coupled with the clock counter **430** (e.g., a finite state machine (FSM), or a respective software stack of the receiver **104** can select a period for detecting the number of pulses. In at least one embodiment, the processing circuitry **132** can select a period from a list of programmed periods. In some embodiments, the processing circuitry **132** can select the period based on reliability constraints of components of the receiver **104**, a transmitter **124**, or the link **220**. For example, if receiver **104** of device **110** is powered down, transmit equalization of the transmitter **124** of device **112** can be disabled within a first period to avoid overstressing the transmitter **124**—e.g., avoid components of the transmitter **124** from exceeding respective threshold voltages and degrading over time. Accordingly, in such embodiments, the processing circuitry **132** can select a period less than the first period associated with the transmitter **124** to ensure the transmit equalization is disabled before degradation occurs. In some embodiments, the processing circuitry **132** can also enable or disable the clock counter based on

transmitting an enable signal **435**.

(37) In at least one embodiment, after detecting the number of pulses in the period, the clock counter **430** can compare the detected number of pulses with an expected number of pulses (e.g., predetermined condition) received via configuration signal **440**. That is, the processing circuitry **132** or respective software stack of device **112** can program the expected number of pulses to the clock counter **430** via configuration signal **440**. In some embodiments, the expected number of pulses is associated with a number of pulses in clock pattern **310**—e.g., associated with a periodic clock signal **305** transmitted during normal operations. In some embodiments, the expected number of pulses is associated with a divided clock pattern **310**. That is, the expected number of pulses can be different based on an amount the incoming clock signal **305** is divided. As the processing circuitry **132** (or the software stack) determines the amount to divide the incoming clock signal **305** by, the processing circuitry **132** (or software stack) can determine and program the expected number of pulses of the divided clock to the clock counter **430** via the configuration signal **440**.

(38) In some embodiments, if the detected number of pulses is different than the expected number of pulses, the clock counter **430** can determine the link **220** at the other device (e.g., device **110**) is not active—e.g., is being shut down or reset. For example, the clock counter **430** can determine the detected number of pulses is different than the expected number of pulses when the device **110** is driving the static clock **315** or not driving a clock signal (e.g., the receiver **104** receives the disabled clock signal **320**) as described with reference to FIG. **3**. In such embodiments, the clock counter **345** can transmit a local link status **325** to the reliability hardware (e.g., reliably hardware **330** as described with reference to FIG. **3**) of device **112**, indicating that the link is not active. Accordingly, the reliability hardware can proceed with initiating and managing a power-down sequence for device **112** as described with reference to FIG. **5**.

(39) In some embodiments, if the detected number of pulses is the same as (e.g., satisfies) the expected number of pulses, the clock counter **430** can determine the link **220** at the other device (e.g., device **110**) is active—e.g., is transmitting data. For example, the clock counter **430** can determine the detected number of pulses that satisfies the expected number of pulses when the device **110** is driving the clock pattern **310** as described with reference to FIG. **3**. In such embodiments, the clock counter **430** can transmit a link status **325** to the reliability hardware **330** of device **112**, indicating the link is active. Accordingly, the reliability hardware can refrain from initiating the power-down sequence for device **112** as described with reference to FIG. **5**.

(40) In some embodiments, the clock counter **430** can be configured to detect the number of pulses in the received clock signal **305** as being the same as the expected number of pulses for multiple periods. For example, the clock counter **430** can be configured to detect the number of pulses as being the same as the expected number of pulses for two (2) periods where each period has the same value. In other embodiments, the clock counter **430** can detect the number of pulses as being the same as the expected number of pulses for greater than two (2) periods. That is, the clock counter **430** can refrain from indicating the link is active or inactive until the specified number of periods has elapsed. In some embodiments, having the specified number of periods be greater than one (1) can increase accuracy and reduce false positives.

(41) FIG. **5** illustrates an example communication system **500** for reliable link management according to at least one embodiment. In at least one embodiment, communication system **500** is an example of communication system **100**, **200**, **300**, or **400** as described with reference to FIGS. **1-4**. The communication system **500** includes a device **110** and a device **112** as described with reference to FIG. **1**. Device **110** and device **112** can be coupled to a link **220** as described with reference to FIG. **2**—e.g., a GRS link **220** having one or more data lanes **202** and a clock lane **205**. Device **110** can include a ramp-down component **505**, a transmit equalization (TX EQ) disabler **510**, receiver phase-locked loop (RX PLL) disabler **515**, disabler **520**, physical layer controls **535**, physical layer **540**-*a*, and ramp-down trigger generator **545**. In an embodiment, the ramp-down component **505**, TX EQ Disabler **510**, RX PLL disabler **515**, disabler **520**, and ramp-down trigger generator **545** can

be considered part of reliability hardware **330** as described with reference to FIG. **3**. Device **112** can include physical layer **540**-*b* and link status detection circuitry **115** as described with reference to FIG. **1**. Although FIG. **5** illustrates different circuitry in device **110** and device **112**, each device can have the same components. For example, device **112** can also include a ramp-down component **505**, a transmit equalization (TX EQ) disabler **510**, receiver phase-locked loop (RX PLL) disabler **515**, disabler **520**, physical layer controls **535**, and ramp-down trigger generator **545**, while device **110** can also include link status detection circuitry **115**. In at least one embodiment, the ramp-down component **505**, TX EQ disabler **510**, RX PLL disabler **515**, and disabler **520** are examples of finite state machines (FSMs).

(42) In at least one embodiment, ramp-down trigger generator **545** is configured to generate a ramp-down trigger **547** based on receiving either a reset signal **550**, software signal **55**, or link status **325**-*a*. For example, the ramp-down trigger generator **545** can generate the ramp-down trigger based on receiving a reset signal **550**. In some embodiments, the ramp-down trigger generator **545** can receive the reset signal **550** from a hardware component. For example, the ramp-down trigger generator **545** can receive the reset signal **550** from error circuitry—e.g., the error circuitry can transmit the reset signal **550** to the ramp-down trigger generator **545** when a number of errors in the communication system **500** or device **110** exceed an error threshold. In some embodiments, the ramp-down trigger generator **545** can receive the reset signal **550** from a software stack associated with the device **110**. In some embodiments, the reset signal **550** can indicate to the ramp-down trigger generator **545** to reset the link **220** or components associated with the link **220** at the device **110**—e.g., reset a receiver **104** or transmitter **124** of device **110**. In some embodiments, the ramp-down trigger generator **545** can generate the ramp-down trigger **547** based on receiving the software signal **555** from the software stack. In some embodiments, the software signal **555** can indicate to the ramp-down trigger generator **545** shut down or power off the link **220** or components associated with the link **220** at the device **110**. In at least one embodiment, the ramp-down trigger generator **545** is configured to generate the ramp-down trigger **547** based on receiving the link status **325**-*a* from the link status detection circuitry **115** of device **110**. For example, the ramp-down trigger generator **545** can generate the ramp-down trigger **547** when the link status **325**-*a* indicates the link is inactive at the device **112** as described with reference to FIGS. **3** and **4**—e.g., device **112** drives the static clock **315**, or the receiver **104** of device **110** receives the disabled clock **320**. In some embodiments, the ramp-down trigger generator **545** can generate the ramp-down trigger **547** based on receiving any combination of the reset signal **550**, software signal **555**, and link status **325**-*a*—e.g., based on receiving reset signal **550** and software signal **555**, or the software signal **555** and link status **325**-*a*, or receiving reset signal **550**, software signal **555**, and link status **325**-*a*, or any combination thereof. In some embodiments, the ramp-down trigger **547** can indicate to the ramp-down component **505** to initiate a power down or ramp-down sequence associated with device **110**.

(43) In at least one embodiment, ramp-down component **505** is configured to receive the ramp-down trigger **547**. In some embodiments, the ramp-down component **505** is configured to initiate the ramp-down sequence associated with device **110** based on receiving the ramp-down trigger **547**. In such embodiments, the ramp-down component **505** can transmit messages, signals, or commands to the physical layer control **535** to shut down or reset analog components in the physical layer **540**-*a* associated with the power-down sequence. In some embodiments, the ramp-down component **505** is configured to transmit messages, commands, or signals to the TX EQ disabler **510**, RX PLL disabler **515**, and disabler **520** to disable their respective components. In at least one embodiment, the software stack can program the ramp-down sequence to the ramp-down component **505**—e.g., the ramp-down component **505** can initiate and manage the ramp-down sequence indicated by the software stack. Accordingly, the ramp-down component **505** can execute the same ramp-down sequence each time the ramp-down trigger **547** is received. In at least one embodiment, the ramp-down component **505** can indicate to the transmitter **124** to transmit a static

clock **315** as described with reference to FIG. **3**—e.g., to alert device **112** that the link **220** and associated components are powering down or resetting at device **110**. In some embodiments, the ramp-down component **505** is configured to ramp down the transmitter **124** before ramping down the receiver **104**. For example, the transmitter **124** can be powered down first as it can otherwise be overstressed or exceed a respective voltage threshold. In other examples, the transmitter **124** can be disabled first to enable time for device **112** to initiate a safe ramp-down sequence and reduce the time that components of device **112** are exposed to unfavorable conditions. For example, as described above, if the receiver **104** of device **110** is disabled while the transmitter **124** of device **112** is enabled, the transmitter **124** of device **112** can undergo additional stress and degrade. Accordingly, device **110** can disable its transmitter **124** first, allowing device **112** to detect the shutdown over the clock lane **205** (e.g., when the clock lane **205** is not driven) and quickly shut down its transmitter **124** before the receiver **104** of device **110** is disabled. In some embodiments, the ramp-down component **505** can ramp down the transmitter **124** at a different time—e.g., the ramp-down component **505** can disable components based on an optimal ramp-down sequence that preserves a reliability of components of device **110**.

(44) In at least one embodiment, the TX EQ disabler **510** is configured to receive the link status **325-**_a_ or a command/signal from the ramp-down component **505**. In at least one embodiment, the TX EQ disabler **510** is configured to disable a transmitter equalization of a transmitter **124** of device **110** based on receiving the link status **325-**_a_ or signal from ramp-down component **505**. For example, the TX EQ disabler **510** can disable the transmitter equalization of transmitter **124** when the link status **325-**_a_ indicates the link **220** is powering down or reset at the device **112**. In some embodiments, the TX EQ disabler **510** can refrain from disabling the transmitter equalization when the link status **325-**_a_ indicates the link **220** is active—e.g., the link status detection circuitry **115** at the device **110** receives the clock pattern **310**.

(45) In at least one embodiment, the RX PLL disabler **515** is configured to receive the link status **325-**_a_ or a command/signal from the ramp-down component **505**.

(46) In at least one embodiment, the disabler **520** is configured to receive the link status **325-**_a_ or a command/signal from the ramp-down component **505**. In at least one embodiment, the disabler **520** is configured to disable other reliability critical components of device **110** based on receiving the link status **325-**_a_ or signal from ramp-down component **505**—e.g., other components that are reset or powered down as quick as possible after determining the link **220** is reset or powered down at the device **112**. For example, the disabler **520** can disable the other reliability components when the link status **325-**_a_ indicates the link **220** is powering down or reset at the device **112**. In some embodiments, the disabler **520** can refrain from disabling the other reliability critical components when the link status **325-**_a_ indicates the link **220** is active—e.g., the link status detection circuitry **115** at the device **110** receives the clock pattern **310**.

(47) It should be noted that TX EQ disabler **510**, RX PLL disabler **515**, and disabler **520** are examples of FSMs the device **110** can utilize to perform the ramp-down sequence. In other embodiments, the device **110** can include additional FSMs for other components that are sensitive to the link **220** being disabled at device **112**. In some embodiments, the TX EQ disabler **510**, RX PLL disabler **515**, and disabler **520** can be part of ramp-down component **505**—e.g., TX EQ disabler **510**, RX PLL disabler **515**, and disabler **520** can be separate when it is critical to disable the TX EQ or RX PLL as quick as possible after detecting the inactive link **220** to avoid reliability issues. In other embodiments, the device **110** can follow different ramp-down procedures. For example, device **110** can perform any combination of the following operations in any order after determining the link **220** is inactive; disable reliability sensitive components of the transmitter **124**, disable reliability sensitive components of the receiver **104**, shut down the transmitter **124** completely, and shut down the receiver **104** completely.

(48) Physical layer control **535** can be an interface between digital logic (e.g., the link status detection circuitry **115**, TX EQ disabler **510**, RX PLL disabler **515**, disabler **520**, ramp-down

component **505**, etc.) and analog components of the physical layer **540**-*a*. In some embodiments, the physical layer control **535** can be an example of flops or registers. In some embodiments, the physical layer control **535** can disable analog circuitry of the physical layer **540**-*a* indicated by the ramp-down component **505**, TX EQ disabler **510**, RX PLL disabler **515**, and disabler **520**.

(49) In at least one embodiment, physical layer **540**-*a* and physical layer **540**-*b* are configured to transmit and receive signals and data over the link **220**. The physical layer **540**-*a* and physical layer **540**-*b* can include the transmitter **124** and receiver **104** as described with reference to FIG. **1**. In some embodiments, the physical layers **540** can be examples of analog circuitry configured to transmit and receive signals over the link **220**.

(50) As described with reference to FIGS. **2**-**4**, the link status detection circuitry **115** is configured to detect a number of pulses of a clock signal received and determine if the detected number of pulses satisfies an expected number of pulses. In some embodiment, the link status detection circuitry **115** can generate a link status **325**-*b* that indicates the link **220** is active when the detected number of pulses satisfies the expected number of pulses. In such embodiments, the device **112** can continue operating normally—e.g., continue transmitting, processing, and receiving data. In some embodiments, the link status detection circuitry **115** can generate a link status **325**-*b* that indicates the link **220** is inactive when the detected number of pulses fails to satisfy the expected number of pulses. In such embodiments, the TX EQ disabler **510**, RX PLL disabler **515**, disabler **520**, and ramp-down component **505** of device **112** can proceed with a power-down or ramp-down sequence of device **112**.

(51) FIG. **6** illustrates an example flow diagram of a method **600** for reliable link management for a high-speed interconnect. The method **600** can be performed by processing logic comprising hardware, software, firmware, or any combination thereof. In at least one embodiment, the method **600** is performed by the transmitter **124** or receiver **104** of the first device **110** or second device **112** as described with reference to FIG. **3**—e.g., by TX EQ disabler **510**, RX PLL disabler **515**, disabler **520**, ramp-down component **505**, ramp-down trigger generator **545**, and link status detection circuitry **115**. Although shown in a particular sequence or order, unless otherwise specified, the order of the processes can be modified. Thus, the illustrated embodiments should be understood only as examples, and the illustrated processes can be performed in a different order, and some processes can be performed in parallel. Additionally, one or more processes can be omitted in various embodiments. Thus, not all processes are required in every embodiment. Other diagrams illustrating a method for reliable link management are possible.

(52) At operation **605**, processing logic can receive a signal via one or more lanes associated with transmitting a clock signal. For example, a system can include a link (e.g., link **220**) including one or more lanes associated with transmitting data (e.g., data lanes **202** and **203**) and one or more lanes associated with transmitting the clock signal (e.g., clock lane **205**). In some embodiments, the processing logic can receive a set of bits associated with a pattern via the one or more lanes associated with transmitting the clock signal, where the pattern is different than a second pattern associated with a data transmission operation—e.g., the signal received can be the set of bits. For example, the processing logic can be in a second device (e.g., device **112**), and a first device (e.g., device **110**) can transmit the set of bits. In some embodiments, the set of bits can be associated with static clock **315**, and the second pattern can be associated with clock pattern **310** as described with reference to FIG. **3**. In some embodiments, the processing logic can receive a second set of bits associated with the second pattern via the one or more lanes associated with transmitting the clock signal—e.g., the processing logic can receive the clock pattern **310**. In at least one embodiment, the processing logic can receive the signal in response to the first device initiating a power-down sequence (e.g., a second power-down sequence). In some embodiments, the processing logic can determine the number of pulses associated with the set of bits fails to satisfy a predetermined condition relating to a specified number of pulses for a period in response to the first or second device transmitting the signal—e.g., fail to satisfy an expected number of pulses for the period as

described with reference to FIGS. **3** and **4**.

(53) At operation **610**, processing logic can determine a number of pulses associated with the signal (e.g., or the set of bits or the second set of bits) over a period. In at least one embodiment, the processing logic is configured to select the period during which the number of pulses is determined from a plurality of periods—e.g., the period is configurable. In some embodiments, the processing logic can select the period based on a reliability condition of a component of the first device or second device—e.g., select the period based on a concern that transmitter **124** of the second device can overstress enabled while the receiver of the first device is disabled. In some embodiments, the processing logic can select a subset of a set of dividers (e.g., dividers **410**) to divide the signal received, where determining the number of pulses associated with the signal over the period is in response to selecting the subset of the set of dividers.

(54) At operation **615**, processing logic can determine the number of pulses associated with the signal fails to satisfy a predetermined condition relating to a specified number of pulses for the period. For example, the processing logic can determine the number of pulses associated with the set of bits fails to satisfy the predetermined condition relating to the specified number of pulses for the period. In some embodiments, the processing logic can determine the number of pulses of the signal received in response to the first device initiating the power-down sequence fails to satisfy the predetermined condition—e.g., the number of pulses detected during the period for the static clock **315** and disabled clock **320** fail to satisfy the predetermined condition relating to the specified number of pulses for the period. In some embodiments, if the processing logic determines the number of pulses fails to satisfy the predetermined condition, the processing logic can proceed to operation **620**. In at least one embodiment, the processing logic can determine the number of pulses associated with the signal (e.g., the second set of bits) satisfy the predetermined condition relating to the specified number of pulses for the period—e.g., the number of pulses detected for the period while receiving the clock pattern **310** satisfies the predetermined condition. In such embodiments, the processing logic can refrain from initiating a power-down sequence in response to determining the number of pulses that satisfies the predetermined condition. In at least one embodiment, the processing logic can determine the number of pulses associated with the signal satisfies the predetermined condition relating to the specified number of pulses for the period a quantity of times, where the quantity of times satisfies a threshold number of times the signal satisfies the number of pulses for the period—e.g., the processing logic can determine the number of pulses over the period fails to satisfy the predetermined for multiple periods.

(55) At operation **620**, the processing logic can initiate a power-down sequence in response to determining the number of pulses that fails to satisfy the predetermined condition relating to the specified number of pulses for the period. For example, the ramp-down trigger generator **545** can generate a ramp-down trigger **547** in response to the number of pulses failing to satisfy the predetermined condition. In at least one embodiment, the processing logic can initiate the power-down sequence in response to receiving a reset signal from a component associated with the link—e.g., based on receiving reset signal **550** from a hardware component associated with the link as described with reference to FIG. **5**. In at least one embodiment, the processing logic can initiate the power-down sequence in response to receiving a second signal from a software stack associated with the device—e.g., based on receiving software signal **555** as described with reference to FIG. **5**.

(56) At operation **625**, the processing logic can execute the power-down sequence. For example, the processing logic can disable transmitter equalization in response to determining the number of pulses does not satisfy the predetermined condition. In some embodiments, the processing logic can disable a receiver in response to disabling the transmitter. In some embodiments, the processing logic can execute a power-down sequence as described with reference to FIG. **5**.

(57) FIG. **7** illustrates a computer system **700** including a transceiver including a chip-to-chip interconnect, in accordance with at least one embodiment. In at least one embodiment, computer system **700** may be a system with interconnected devices and components, an SOC, or some

combination. In at least one embodiment, computer system **700** is formed with a processor **702** that may include execution units to execute an instruction. In at least one embodiment, computer system **700** may include, without limitation, a component, such as processor **702**, to employ execution units including logic to perform algorithms for processing data. In at least one embodiment, computer system **700** may include processors, such as PENTIUM® Processor family, Xeon™, Itanium®, XScale™ and/or StrongARM™, Intel® Core™, or Intel® Nervana™ microprocessors available from Intel Corporation of Santa Clara, California, although other systems (including PCs having other microprocessors, engineering workstations, set-top boxes and like) may also be used. In at least one embodiment, computer system **700** may execute a version of WINDOWS' operating system available from Microsoft Corporation of Redmond, Wash., although other operating systems (UNIX and Linux for example), embedded software, and/or graphical user interfaces may also be used.

(58) In at least one embodiment, computer system **700** may be used in other devices such as handheld devices and embedded applications. Some examples of handheld devices include cellular phones, Internet Protocol devices, digital cameras, personal digital assistants ("PDAs"), and handheld PCs. In at least one embodiment, embedded applications may include a microcontroller, a digital signal processor (DSP), an SoC, network computers ("NetPCs"), set-top boxes, network hubs, wide area network ("WAN") switches, or any other system that may perform one or more instructions. In an embodiment, computer system **700** may be used in devices such as graphics processing units (GPUs), network adapters, central processing units and network devices such as switch (e.g., a high-speed direct GPU-to-GPU interconnect such as the NVIDIA GH100 NVLINK or the NVIDIA Quantum 2 64 Ports InfiniBand NDR Switch).

(59) In at least one embodiment, computer system **700** may include, without limitation, processor **702** that may include, without limitation, one or more execution units **707** that may be configured to execute a Compute Unified Device Architecture ("CUDA") (CUDA® is developed by NVIDIA Corporation of Santa Clara, CA) program. In at least one embodiment, a CUDA program is at least a portion of a software application written in a CUDA programming language. In at least one embodiment, computer system **700** is a single-processor desktop or server system. In at least one embodiment, computer system **700** may be a multiprocessor system. In at least one embodiment, processor **702** may include, without limitation, a CISC microprocessor, a RISC microprocessor, a VLIW microprocessor, a processor implementing a combination of instruction sets, or any other processor device, such as a digital signal processor, for example. In at least one embodiment, processor **702** may be coupled to a processor bus **710** that may transmit data signals between processor **702** and other components in computer system **700**.

(60) In at least one embodiment, processor **702** may include, without limitation, a Level 1 ("L1") internal cache memory ("cache") **704**. In at least one embodiment, processor **702** may have a single internal cache or multiple levels of internal cache. In at least one embodiment, cache memory may reside external to processor **702**. In at least one embodiment, processor **702** may also include a combination of both internal and external caches. In at least one embodiment, a register file **706** may store different types of data in various registers including, without limitation, integer registers, floating point registers, status registers, and instruction pointer register.

(61) In at least one embodiment, execution unit **707**, including, without limitation, logic to perform integer and floating point operations, also resides in processor **702**. Processor **702** may also include a microcode ("ucode") read only memory ("ROM") that stores microcode for certain macro instructions. In at least one embodiment, execution unit **707** of processor **702** may include logic to handle a packed instruction set **709**. In at least one embodiment, by including packed instruction set **709** in an instruction set of a general-purpose processor **702**, along with associated circuitry to execute instructions, operations used by many multimedia applications may be performed using packed data in a general-purpose processor **702**. In at least one embodiment, many multimedia applications may be accelerated and executed more efficiently by using full width of a processor's

data bus for performing operations on packed data, which may eliminate a need to transfer smaller units of data across a processor's data bus to perform one or more operations one data element at a time.

(62) In at least one embodiment, an execution unit may also be used in microcontrollers, embedded processors, graphics devices, DSPs, and other types of logic circuits. In at least one embodiment, computer system **700** may include, without limitation, a memory **720**. In at least one embodiment, memory **720** may be implemented as a DRAM device, an SRAM device, flash memory device, or other memory device. Memory **720** may store instruction(s) **719** and/or data **721** represented by data signals that may be executed by processor **702**.

(63) In at least one embodiment, a system logic chip may be coupled to processor bus **710** and memory **720**. In at least one embodiment, the system logic chip may include, without limitation, a memory controller hub ("MCH") **716**, and processor **702** may communicate with MCH **716** via processor bus **710**. In at least one embodiment, MCH **716** may provide a high bandwidth memory path **718** to memory **720** for instruction and data storage and for storage of graphics commands, data and textures. In at least one embodiment, MCH **716** may direct data signals between processor **702**, memory **720**, and other components in computer system **700** and to bridge data signals between processor bus **710**, memory **720**, and a system I/O **722**. In at least one embodiment, a system logic chip may provide a graphics port for coupling to a graphics controller. In at least one embodiment, MCH **716** may be coupled to memory **720** through high bandwidth memory path **718**, and graphics/video card **712** may be coupled to MCH **716** through an Accelerated Graphics Port ("AGP") interconnect **714**.

(64) In at least one embodiment, computer system **700** may use system I/O **722** that is a proprietary hub interface bus to couple MCH **716** to I/O controller hub ("ICH") **730**. In at least one embodiment, ICH **730** may provide direct connections to some I/O devices via a local I/O bus. In at least one embodiment, a local I/O bus may include, without limitation, a high-speed I/O bus for connecting peripherals to memory **720**, a chipset, and processor **702**. Examples may include, without limitation, an audio controller **729**, a firmware hub ("flash BIOS") **728**, a transceiver **726**, a data storage **724**, a legacy I/O controller **723** containing a user input interface **725** and a keyboard interface, a serial expansion port **727**, such as a USB, and a network controller **734**. Data storage **724** may comprise a hard disk drive, a floppy disk drive, a CD-ROM device, a flash memory device, or other mass storage device.

(65) In at least one embodiment, FIG. **7** illustrates a system, which includes interconnected hardware devices or "chips" in the transceiver **726**—e.g., the transceiver **726** includes a chip-to-chip interconnect including the first device **110** and second device **112** as described with reference to FIG. **1**). In at least one embodiment, FIG. **7** may illustrate an exemplary SoC. In at least one embodiment, devices illustrated in FIG. **7** may be interconnected with proprietary interconnects, standardized interconnects (e.g., PCIe), or some combination thereof and utilize a GRS link **220** as described with reference to FIG. **2**. In at least one embodiment, one or more components of system **700** are interconnected using compute express link ("CXL") interconnects. In an embodiment, the transceiver **726** can include link status detection circuitry **115**. In some embodiments, link status detection circuitry **115** can determine whether the link is powered on at the other device. If the link status detection circuitry **115** determines the link is not powered on at the other device, the link status detection circuitry **115** can generate an indication (e.g., a link status) that the link is powered down at the other device and transmit the indication to a ramp-down component as described with reference to FIG. **5**. The ramp-down component can initiate a power-down sequence in response to receiving the indication, enabling the chip to follow the power-down sequence and maintain the reliability of the link components and components associated with the link. Accordingly, the link status detection circuitry **115** can be configured for reliable link management as described with reference to FIGS. **2-6**.

(66) Other variations are within spirit of present disclosure. Thus, while disclosed techniques are

susceptible to various modifications and alternative constructions, certain illustrated embodiments thereof are shown in drawings and have been described above in detail. It should be understood, however, that there is no intention to limit the disclosure to a specific form or forms disclosed, but on the contrary, the intention is to cover all modifications, alternative constructions, and equivalents falling within the spirit and scope of the disclosure, as defined in appended claims.

(67) Use of terms "a" and "an" and "the" and similar referents in the context of describing disclosed embodiments (especially in the context of following claims) are to be construed to cover both singular and plural, unless otherwise indicated herein or clearly contradicted by context, and not as a definition of a term. Terms "comprising," "having," "including," and "containing" are to be construed as open-ended terms (meaning "including, but not limited to,") unless otherwise noted. "Connected," when unmodified and referring to physical connections, is to be construed as partly or wholly contained within, attached to, or joined together, even if there is something intervening. Recitations of ranges of values herein are merely intended to serve as a shorthand method of referring individually to each separate value falling within the range, unless otherwise indicated herein, and each separate value is incorporated into the specification as if it were individually recited herein. In at least one embodiment, the use of the term "set" (e.g., "a set of items") or "subset" unless otherwise noted or contradicted by context, is to be construed as a nonempty collection comprising one or more members. Further, unless otherwise noted or contradicted by context, the term "subset" of a corresponding set does not necessarily denote a proper subset of the corresponding set, but subset and corresponding set may be equal.

(68) Conjunctive language, such as phrases of the form "at least one of A, B, and C," or "at least one of A, B and C," unless specifically stated otherwise or otherwise clearly contradicted by context, is otherwise understood with the context as used in general to present that an item, term, etc., may be either A or B or C, or any nonempty subset of the set of A and B and C. For instance, in an illustrative example of a set having three members, conjunctive phrases "at least one of A, B, and C" and "at least one of A, B and C" refer to any of the following sets: {A}, {B}, {C}, {A, B}, {A, C}, {B, C}, {A, B, C}. Thus, such conjunctive language is not generally intended to imply that certain embodiments require at least one of A, at least one of B and at least one of C each to be present. In addition, unless otherwise noted or contradicted by context, the term "plurality" indicates a state of being plural (e.g., "a plurality of items" indicates multiple items). In at least one embodiment, the number of items in a plurality is at least two, but can be more when so indicated either explicitly or by context. Further, unless stated otherwise or otherwise clear from context, the phrase "based on" means "based at least in part on" and not "based solely on."

(69) Operations of processes described herein can be performed in any suitable order unless otherwise indicated herein or otherwise clearly contradicted by context. In at least one embodiment, a process such as those processes described herein (or variations and/or combinations thereof) is performed under control of one or more computer systems configured with executable instructions and is implemented as code (e.g., executable instructions, one or more computer programs or one or more applications) executing collectively on one or more processors, by hardware or combinations thereof. In at least one embodiment, code is stored on a computer-readable storage medium, for example, in the form of a computer program comprising a plurality of instructions executable by one or more processors. In at least one embodiment, a computer-readable storage medium is a non-transitory computer-readable storage medium that excludes transitory signals (e.g., a propagating transient electric or electromagnetic transmission) but includes non-transitory data storage circuitry (e.g., buffers, cache, and queues) within transceivers of transitory signals. In at least one embodiment, code (e.g., executable code or source code) is stored on a set of one or more non-transitory computer-readable storage media having stored thereon executable instructions (or other memory to store executable instructions) that, when executed (i.e., as a result of being executed) by one or more processors of a computer system, cause a computer system to perform operations described herein. In at least one embodiment, a set of

non-transitory computer-readable storage media comprises multiple non-transitory computer-readable storage media and one or more of individual non-transitory storage media of multiple non-transitory computer-readable storage media lack all of the code while multiple non-transitory computer-readable storage media collectively store all of the code. In at least one embodiment, executable instructions are executed such that different instructions are executed by different processors.

(70) Accordingly, in at least one embodiment, computer systems are configured to implement one or more services that singly or collectively perform operations of processes described herein and such computer systems are configured with applicable hardware and/or software that enable the performance of operations. Further, a computer system that implements at least one embodiment of present disclosure is a single device and, in another embodiment, is a distributed computer system comprising multiple devices that operate differently such that distributed computer system performs operations described herein and such that a single device does not perform all operations.

(71) Use of any and all examples, or exemplary language (e.g., "such as") provided herein, is intended merely to better illuminate embodiments of the disclosure and does not pose a limitation on the scope of the disclosure unless otherwise claimed. No language in the specification should be construed as indicating any non-claimed element as essential to the practice of the disclosure.

(72) All references, including publications, patent applications, and patents, cited herein are hereby incorporated by reference to the same extent as if each reference were individually and specifically indicated to be incorporated by reference and were set forth in its entirety herein.

(73) In description and claims, terms "coupled" and "connected," along with their derivatives, may be used. It should be understood that these terms may not be intended as synonyms for each other. Rather, in particular examples, "connected" or "coupled" may be used to indicate that two or more elements are in direct or indirect physical or electrical contact with each other. "Coupled" may also mean that two or more elements are not in direct contact with each other, but yet still co-operate or interact with each other.

(74) Unless specifically stated otherwise, it may be appreciated that throughout specification terms such as "processing," "computing," "calculating," "determining," or like, refer to action and/or processes of a computer or computing system, or similar electronic computing device, that manipulate and/or transform data represented as physical, such as electronic, quantities within computing system's registers and/or memories into other data similarly represented as physical quantities within computing system's memories, registers or other such information storage, transmission or display devices.

(75) In a similar manner, the term "processor" may refer to any device or portion of a device that processes electronic data from registers and/or memory and transform that electronic data into other electronic data that may be stored in registers and/or memory. A "computing platform" may comprise one or more processors. As used herein, "software" processes may include, for example, software and/or hardware entities that perform work over time, such as tasks, threads, and intelligent agents. Also, each process may refer to multiple processes, for carrying out instructions in sequence or in parallel, continuously or intermittently. In at least one embodiment, terms "system" and "method" are used herein interchangeably insofar as the system may embody one or more methods and methods may be considered a system.

(76) In the present document, references may be made to obtaining, acquiring, receiving, or inputting analog or digital data into a subsystem, computer system, or computer-implemented machine. In at least one embodiment, the process of obtaining, acquiring, receiving, or inputting analog and digital data can be accomplished in a variety of ways such as by receiving data as a parameter of a function call or a call to an application programming interface. In at least one embodiment, processes of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a serial or parallel interface. In at least one embodiment, processes of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished

by transferring data via a computer network from providing entity to acquiring entity. In at least one embodiment, references may also be made to providing, outputting, transmitting, sending, or presenting analog or digital data. In various examples, processes of providing, outputting, transmitting, sending, or presenting analog or digital data can be accomplished by transferring data as an input or output parameter of a function call, a parameter of an application programming interface or inter-process communication mechanism.

(77) Although descriptions herein set forth example embodiments of described techniques, other architectures may be used to implement described functionality, and are intended to be within the scope of this disclosure. Furthermore, although specific distributions of responsibilities may be defined above for purposes of description, various functions and responsibilities might be distributed and divided in different ways, depending on circumstances.

(78) Furthermore, although the subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that subject matter claimed in appended claims is not necessarily limited to specific features or acts described. Rather, specific features and acts are disclosed as exemplary forms of implementing the claims.

Claims

1. A device comprising: receiver circuitry to receive incoming signals on a clock lane and data lanes, wherein the clock lane and data lanes are part of a Ground Reference Signaling (GRS) link, wherein the incoming signals on the clock lane comprise a forwarded clock associated with data signals on the data lanes of the GRS link, wherein the GRS link comprises a clock pattern on the clock lane, the clock pattern being associated with communicating data on the data lanes; and detection circuitry to: monitor the incoming signals on the clock lane; determine that an incoming pattern of the incoming signals on the clock lane does not correspond to the clock pattern associated with communicating data on the data lanes; and initiate a power-down sequence in response to determining that the incoming pattern does not correspond to the clock pattern.

2. The device of claim 1, wherein the clock lane and data lanes are part of the GRS link between the device and a second device, wherein the forwarded clock is an indication of whether the GRS link is active, wherein the GRS link is active responsive to the forwarded clock having the clock pattern, wherein the GRS link is inactive responsive to the forwarded clock not having the clock pattern.

3. The device of claim 1, wherein: the detection circuitry, to determine that the incoming pattern of the incoming signals on the clock lane does not correspond to the clock pattern, is to: determine a number of pulses associated with the incoming signals on the clock lane over a period; determine the number of pulses fail to satisfy a predetermined condition relating to a specified number of pulses for the period; and the detection circuitry is to initiate the power-down sequence in response to determining that the number of pulses fails to satisfy the predetermined condition.

4. The device of claim 3, wherein the detection circuitry is further to refrain from initiating the power-down sequence in response to determining the number of pulses satisfies the predetermined condition.

5. The device of claim 1, wherein the incoming signals on the clock lane comprise a forwarded clock, wherein the detection circuitry comprises: link status detection circuitry to receive the forwarded clock and output a link status; and reliability hardware to receive the link status from the link status detection circuitry, wherein the reliability hardware is to selectively disable circuits of the receiver circuitry or perform a controlled shut down of the receiver circuitry in response to the link status indicating an inactive state.

6. The device of claim 5, wherein the reliability hardware comprises at least one of: a ramp-down trigger generator to generate a trigger for a ramp-down component in response to the link status indicating an inactive state; the ramp-down component to provide at least one of a message, a

command, or a signal to initiate the power-down sequence in response to the trigger from the ramp-down trigger generator; and a disabler to disable one or more components of the receiver circuitry in response to receiving the at least one of the message, the command, or the signal from the ramp-down component.

7. The device of claim 5, wherein the link status detection circuitry comprises: a sampler to sample the forwarded clock on the clock lane; a clock divider coupled to the sampler, the clock divider to generate a divided clock that is slower than the forwarded clock; a low pass filter coupled to the clock divider, the low pass filter to remove short-term fluctuations in the divided clock; and a counter coupled to the low pass filter, the counter to count a number of pulses in the divided clock during a detection window, wherein the forwarded clock is active in response to the counter having a deterministic count value in the detection window.

8. The device of claim 1, further comprising: transmitter circuitry coupled to the detection circuitry, wherein the detection circuitry is to: disable at least a portion of the transmitter circuitry in response to determining that the incoming pattern does not correspond to the clock pattern; and disable at least a portion of the receiver circuitry in response to disabling at least the portion of the transmitter circuitry.

9. The device of claim 1, wherein, to determine that the incoming pattern of the incoming signals on the clock lane does not correspond to the clock pattern, the detection circuitry is to identify a static pattern in the forwarded clock.

10. The device of claim 1, wherein, to determine that the incoming pattern of the incoming signals on the clock lane does not correspond to the clock pattern, the detection circuitry is to determine that the forwarded clock is no longer driven on the clock lane.

11. A device comprising: receiver circuitry to receive incoming signals on a clock lane and data lanes; and detection circuitry to: monitor the incoming signals on the clock lane; determine that an incoming pattern of the incoming signals on the clock lane does not correspond to a clock pattern associated with communicating data on the data lanes; and initiate a power-down sequence in response to determining that the incoming pattern does not correspond to the clock pattern, wherein the incoming signals on the clock lane comprise a forwarded clock, and wherein, to determine that the incoming pattern of the incoming signals on the clock lane does not correspond to the clock pattern, the detection circuitry is to identify a static pattern in the forwarded clock.

12. The device of claim 11, wherein the clock lane and the data lanes are part of a Ground Reference Signaling (GRS) link, wherein the incoming signals on the data lanes comprise data signals, wherein the forwarded clock is associated with the data signals of the GRS link, wherein the GRS link comprises the clock pattern on the clock lane during active operation.

13. A device comprising: receiver circuitry to receive incoming signals on a clock lane and data lanes; and detection circuitry to: monitor the incoming signals on the clock lane; determine that an incoming pattern of the incoming signals on the clock lane does not correspond to a clock pattern associated with communicating data on the data lanes; and initiate a power-down sequence in response to determining that the incoming pattern does not correspond to the clock pattern, wherein the incoming signals on the clock lane comprise a forwarded clock, and wherein, to determine that the incoming pattern of the incoming signals on the clock lane does not correspond to the clock pattern, the detection circuitry is to determine that the forwarded clock is no longer driven on the clock lane.

14. The device of claim 13, wherein the clock lane and the data lanes are part of a Ground Reference Signaling (GRS) link, wherein the incoming signals on the data lanes comprise data signals, wherein the forwarded clock is associated with the data signals of the GRS link, wherein the GRS link comprises the clock pattern on the clock lane during active operation.

15. A method of operating a device, the method comprising: receiving incoming signals on a clock lane and data lanes; monitoring the incoming signals on the clock lane; determining that an incoming pattern of the incoming signals on the clock lane does not correspond to a clock pattern

associated with communicating data on the data lanes; and initiating a power-down sequence in response to determining that the incoming pattern does not correspond to the clock pattern, wherein the incoming signals on the clock lane comprise a forwarded clock, and wherein determining that the incoming pattern of the incoming signals on the clock lane does not correspond to the clock pattern comprises either: identifying a static pattern in the forwarded clock; or determining that the forwarded clock is no longer driven on the clock lane.

16. The method of claim 15, wherein the incoming signals on the clock lane comprise a forwarded clock, and wherein determining that the incoming pattern of the incoming signals on the clock lane does not correspond to the clock pattern comprises: sampling the forwarded clock on the clock lane; generating a divided clock from the forwarded clock, wherein the divided clock is slower than the forwarded clock; filtering the divided clock to remove short-term fluctuations in the divided clock; and counting, using a counter, a number of pulses in the divided clock during a detection window, wherein the forwarded clock is active in response to the counter having a deterministic count value in the detection window.

17. The method of claim 15, further comprising: disabling at least a portion of transmitter circuitry of the device in response to determining that the incoming pattern does not correspond to the clock pattern; and disabling at least a portion of receiver circuitry of the device in response to disabling at least the portion of the transmitter circuitry.

18. A system comprising: a Ground Reference Signaling (GRS) link comprising a clock lane and data lanes; a first device coupled to the GRS link; and a second device coupled to the GRS link, wherein the second device comprises: receiver circuitry to receive incoming signals on the clock lane and the data lanes, wherein the incoming signals on the clock lane comprise a forwarded clock associated with data signals on the data lanes of the GRS link, wherein the GRS link comprises a clock pattern on the clock lane, the clock pattern being associated with communicating data on the data lanes; and detection circuitry to: monitor the incoming signals on the clock lane; determine that an incoming pattern of the incoming signals on the clock lane does not correspond to the clock pattern associated with communicating data on the data lanes; and initiate a power-down sequence in response to determining that the incoming pattern does not correspond to the clock pattern.

19. The system of claim 18, wherein, to determine that the incoming pattern of the incoming signals on the clock lane does not correspond to the clock pattern, the detection circuitry is to identify a static pattern in the forwarded clock.

20. The system of claim 18, wherein, to determine that the incoming pattern of the incoming signals on the clock lane does not correspond to the clock pattern, the detection circuitry is to determine that the forwarded clock is no longer driven on the clock lane.