

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12395338
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Kim; Hyun et al.

Electronic device for performing token pruning in frequency domain and method for operating the same

Abstract

According to various embodiment of the present disclosure, an electronic device for performing token pruning in a frequency domain may include a processor, and the processor may configured to divide an image frame into a plurality of patches, convert tokens based on the plurality of patches from a spatial domain to a frequency domain through a fast Fourier transform (FFT)-based frequency domain conversion algorithm, the tokens being output from a predetermined transformer block among a plurality of transformer blocks, after performing patch embedding on the plurality of patches, and convert the remaining tokens other than specific tokens from the frequency domain to the spatial domain, after pruning specific tokens from the tokens based on frequency information of the tokens converted into the frequency domain. Various other embodiments are also possible.

Inventors: Kim; Hyun (Seoul, KR), Lee; Jong Ho (Seoul, KR)

Applicant: Foundation for Research and Business, Seoul National University of Science and Technology (Seoul, KR)

Family ID: 1000008762563

Assignee: Foundation for Research and Business, Seoul National University of Science and Technology (Seoul, KR)

Appl. No.: 18/166828

Filed: February 09, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240214205 A1	Jun. 27, 2024

Foreign Application Priority Data

KR	10-2022-0184327	Dec. 26, 2022
----	-----------------	---------------

Publication Classification

Int. Cl.: **G06T5/00** (20240101); **G06T5/10** (20060101); **G06T7/11** (20170101); **G06V10/26** (20220101); **H04L9/32** (20060101)

U.S. Cl.:

CPC **H04L9/3213** (20130101); **G06T5/10** (20130101); **G06T7/11** (20170101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2003/0063809	12/2002	Andrew	382/240	G06T 9/007
2003/0142847	12/2002	Rhoads	704/E19.009	G07F 7/12
2004/0170332	12/2003	Andrew	382/233	H04N 19/36
2013/0145239	12/2012	Pi	714/807	H04L 25/0224
2019/0362477	12/2018	Rahmati	N/A	G06T 5/70
2022/0012848	12/2021	Ranftl	N/A	G06T 3/4046
2022/0094713	12/2021	Lee	N/A	G06F 18/214
2022/0114709	12/2021	Rahmati	N/A	G06T 5/20
2022/0277218	12/2021	Fan	N/A	G06N 3/08

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
115240003	12/2021	CN	N/A

Primary Examiner: Robinson; Christopher B

Attorney, Agent or Firm: Studebaker Brackett PLLC

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS

(1) This application claims priority to and the benefit of Korean Patent Application No. 10-2022-0184327 filed in the Korean Intellectual Property Office on Dec. 26, 2022, the entire contents of which are incorporated herein by reference.

TECHNICAL FIELD

(2) The present disclosure relates to an electronic device for performing token pruning in a frequency domain and a method for operating the same.

BACKGROUND ART

(3) Recently, vision transformer (ViT) family deep learning models have achieved remarkable results in computer vision tasks such as image classification, object detection, image segmentation,

and depth estimation. Based on such development, the ViTs are being studied to be actively used in various applications. For example, the ViTs are widely studied for the purpose of processing images in mobile devices, autonomous vehicles, and drones as well as super-large deep learning models, and a number of domestic and foreign companies, including large companies, mid-sized and small-sized companies, and startups, have much interest in the ViTs. Specifically, in order to efficiently utilize ViT-based applications in devices under embedded environments operating based on a battery, it is essential to reduce the computation amount and power consumption through efficient lightweight technology in consideration of network accuracy and a trade-off of computation amount.

(4) It is known that ViT performance increases in proportion to an increase in parameters and the computation amount. However, because the increase in parameters and the computation amount inevitably results in severe power consumption, it may be considerably difficult to efficiently utilize the ViTs in mobile devices or autonomous vehicles with limited hardware resources. Many studies have been conducted to compress a network in order to solve such a problem and effectively use the ViTs in various applications. One of the most widely known methods, a pruning scheme, is meant to remove relatively less important tokens, layers, and channels from neural networks.

SUMMARY OF THE INVENTION

(5) An aspect of the present disclosure is to provide a method for pruning tokens of a transformer-based neural network that can reduce the computation amount and improve inference speed by effectively removing tokens with unnecessary high frequency components through frequency conversion.

(6) According to various embodiments of the present disclosure, an electronic device for performing token pruning in a frequency domain may include a processor, and the processor may be configured to divide an image frame into a plurality of patches, convert tokens based on the plurality of patches from a spatial domain to a frequency domain through a fast Fourier transform (FFT)-based frequency domain conversion algorithm, the tokens being output from a predetermined transformer block among a plurality of transformer blocks, after performing patch embedding on the plurality of patches, and convert the remaining tokens other than specific tokens from the frequency domain to the spatial domain, after pruning specific tokens from the tokens based on frequency information of the tokens converted into the frequency domain.

(7) According to various embodiments of the present disclosure, a method of operating an electronic device for performing token pruning in a frequency domain may include: dividing an image frame into a plurality of patches using a processor of the electronic device; converting tokens based on the plurality of patches from a spatial domain to a frequency domain through a fast Fourier transform (FFT)-based frequency domain conversion algorithm, the tokens being output from a predetermined transformer block among a plurality of transformer blocks, after performing patch embedding on the plurality of patches, using the processor; and converting the remaining tokens other than specific tokens from the frequency domain to the spatial domain, after pruning specific tokens from the tokens based on frequency information of the tokens converted into the frequency domain, using the processor.

(8) In the present disclosure, high frequency components known to be relatively less important may be removed from an image, and learning may be performed using only low frequency components, thereby not only having a less decrease in accuracy due to lightness than a conventional method but also reducing the learning time due to a decrease in the computation amount and reducing a large amount of CO₂ generated by a GPU.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1 illustrates a block diagram of an electronic device and a network according to various embodiments of the present disclosure.

(2) FIG. 2 is a flowchart illustrating a method for performing token pruning by the electronic device according to various embodiments of the present disclosure.

(3) FIG. 3 is a view illustrating an embodiment in which the electronic device **101** performs the token pruning according to various embodiments of the present disclosure.

DETAILED DESCRIPTION

(4) Hereinafter, embodiments of the present disclosure will be described with reference to the attached drawings. It should be understood that embodiments and the terms used herein are not intended to limit the techniques described herein to specific embodiments, and the invention is to cover all modifications, equivalents, and alternatives falling within the spirit and scope of the invention. Regarding the drawings, like numbers refer to like elements throughout the description of the figures. The singular also includes the plural unless specifically stated otherwise in the phrase. In the present disclosure, the expression “A or B,” “at least one of A or/and B,” or “one or more of A of/and B” may include all possible combinations of listed items. The expressions, such as “first” and “second”, used in the present disclosure can be used to describe various elements without regarding to sequence and/or importance and do not limit corresponding elements but are used only to classify a certain element from another element. When it is described that a certain element (e.g., a first element) is “(operatively or communicatively) coupled with/to” or “connected to” another element (e.g., a second element), it should be understood that the certain element may be connected to another element directly or via another element (e.g., a third element). The expressions “configured to” used in the present disclosure may be interchangeably replaced with, for example, “suitable for,” “having the capacity to,” “designed to,” “adapted to,” “made to,” or “capable of” in accordance with circumstances. In some cases, the expression “device configured to” may indicate the device “capable of” together with other devices or parts. For example, the wording “processor configured to perform A, B, and C” may indicate an exclusive processor (e.g., an embedded processor) for performing corresponding operations or a generic-purpose processor (e.g., a central processing unit (CPU) or an application processor) capable of performing corresponding operations by executing one or more software programs stored in a memory device.

(5) An electric device according to various embodiments of the present disclosure may include, for example, at least one of a smartphone, a tablet PC, a desktop PC, a laptop PC, a netbook computer, a workstation, and a server.

(6) Referring to FIG. 1, the electronic device **101** in a network environment **100** is described in various embodiments. The electronic device **101** may include a bus **110**, a processor **120**, a memory **130**, an input/output interface **150**, a display **160**, and a communication interface **170**. In some embodiments, the electronic device **101** may omit at least one of the components or additionally include another component. The bus **110** may include a circuit that connects the components **110-170** to each other and transmits communication (e.g., control messages or data) between the components. The processor **120** may include one or more of a central processing unit, an application processor, or a communication processor (CP). For example, the processor **120** may perform arithmetic operations or data processing regarding control and/or communication of at least one different component of the electronic device **101**.

(7) The memory **130** may include volatile and/or nonvolatile memory. The memory **130** may store, for example, instructions or data related to at least one different component of the electronic device **101**. According to an embodiment, the memory **130** may store software and/or a program **140**. The program **140** may include, for example, a kernel **141**, middleware **143**, an application programming interface (API) **145** and/or an application program (or “application”) **147**. At least some of the kernel **141**, the middleware **143**, or the API **145** may be referred to as an operating system. The

kernel **141** may control or manage system resources (e.g., the bus **110**, the processor **120**, or the memory **130**) used to execute operations or functions implemented in other programs (e.g., the middleware **143**, the API **145**, or the application program **147**). In addition, the kernel **141** may provide an interface for controlling or managing system resources by accessing individual components of the electronic device **101** in the middleware **143**, the API **145**, or the application program **147**.

(8) The middleware **143** may serve as an intermediary that enables, for example, the API **145** or the application program **147** to communicate with kernel **141** so as to exchange data. In addition, the middleware **143** may process one or more task requests received from the application program **147** according to their priority. For example, the middleware **143** may assign a priority to at least one of application programs **147** to use system resources (e.g., the bus **110**, the processor **120**, the memory **130**, or the like) of the electronic device **101** and may process the one or more task requests. The API **145** is an interface for allowing the application **147** to control functions provided by the kernel **141** or middleware **143**, and may include at least one interface or function (e.g., an instruction) for file control, window control, image processing, or character control. The input/output interface **150** may transmit, for example, instructions or data input from the user or other external devices to other component(s) of the electronic device **101**, or output instructions or data received from other component(s) of the electronic device **101** to the user or other external devices.

(9) The display **160** may include, for example, a liquid crystal display (LCD), a light emitting diode (LED) display, an organic light emitting diode (OLED) display, a microelectromechanical system (MEMS) display, or an electronic paper display. For example, the display **160** may display various contents (e.g., text, images, videos, icons, and/or symbols) to a user. The display **160** may include a touch screen, and may receive, for example, a touch, a gesture, proximity, or hovering input using an electronic pen or a portion of the user's body. The communication interface **170** may establish communication between, for example, the electronic device **101** and external devices (e.g., a first external electronic device **102**, a second external electronic device **104**, or a server **106**). For example, the communication interface **170** may be connected to a network **162** through wireless communication or wired communication to communicate with the external device (e.g., a second external electronic device **104** or a server **106**).

(10) The wireless communication may cellular communication using at least one of, for example, LTE, LTE advance (LTE-A), code division multiple access (CDMA), wideband CDMA (WCDMA), a universal mobile telecommunications system (UMTS), wireless broadband (WiBro), or a global system for mobile communications (GSM). According to an embodiment, the wireless communication may include at least one of, for example, wireless fiber (Wi-Fi), Bluetooth, Bluetooth low power (BLE), Zigbee, near field communication (NFC), magnetic secure transmission, radio frequency (RF), or a body area network (BAN). According to an embodiment, the wireless communication may include a GNSS. The GNSS may be, for example a global positioning system (GPS), a global navigation satellite system (Glonass), a Beidou navigation satellite system (hereinafter referred to as "Beidou") or Galileo, a European global satellite-based navigation system. Hereinafter, in the present disclosure, the "GPS" may be used interchangeably with the "GNSS." The wired communication may include, for example, at least one of universal serial bus (USB), a high-definition multimedia interface (HDMI), a recommended standard 232 (RS-232), power line communication, or a plain old telephone service (POTS). The network **162** may include, for example, at least one of a telecommunication network, a computer network (e.g., a LAN or WAN), the Internet, or a telephony network.

(11) Each of the first and second external electronic devices **102** and **104** may be the same as or different from the electronic device **101**. According to various embodiments, all or some of the operations executed in the electronic device **101** may be executed by another electronic device or a plurality of electronic devices (e.g., the electronic devices **102** and **104**), or the server **106**.

According to one embodiment, when the electronic device **101** needs to perform a function or a

service automatically or on request, the electronic device **101** may request at least some of the associated functions from another device (e.g., the electronic device **102** or **104**, or the server **106**, additionally or instead of autonomously executing the function or the service. The electronic device (e.g., the electronic device **102** or **104**, or the server **106**) may execute the requested function or additional function and transmit results thereof to the electronic device **101**. The electronic device **101** may provide the requested function or service by processing the received results intactly or additionally. To this end, for example, cloud computing, distributed computing, or client-server computing technology may be used.

(12) FIG. **2** is a flowchart illustrating a method for performing the token pruning by the electronic device (e.g., the electronic device **101** of FIG. **1**) according to various embodiments of the present disclosure.

(13) FIG. **3** is a view illustrating an embodiment in which the electronic device **101** performs the token pruning according to various embodiments of the present disclosure.

(14) In operation **201**, according to various embodiments, the electronic device (e.g., the processor **120** of FIG. **1**) may divide an image frame into a plurality of patches. According to an embodiment, the electronic device **101** may obtain the image frame from an internal device (e.g., a camera module) or an external device (e.g., an external server), and may divide the obtained image frame into the plurality of patches. For example, referring to FIG. **3**, the electronic device **101** may divide one image frame **301** into a plurality of patches **302**.

(15) In operation **203**, according to various embodiments, the electronic device (e.g., processor **120** in FIG. **1**) may perform patch embedding on the plurality of patches, and then convert tokens based on the plurality of patches from a spatial domain to a frequency domain through a fast Fourier transform (FFT)-based frequency domain conversion algorithm, the tokens being output from a predetermined transformer block among a plurality of transformer blocks.

(16) According to an embodiment, the electronic device **101** may perform the patch embedding on the plurality of patches **302**. For example, referring to FIG. **3**, the electronic device **101** may perform the patch embedding **301** on a plurality of divided patches **302**. According to an embodiment, the electronic device **101** may perform the patch embedding without generating a CLS token while performing the patch embedding.

(17) According to an embodiment, the electronic device **101** may identify the tokens output from the predetermined transformer block among the plurality of transformer blocks. In this case, the tokens may be generated based on the plurality of patches. In an embodiment, the image token may refer to each patch of the plurality of patches **302**. A transformer is a neural network model referred to as a foundation model that can learn context and meaning by tracking relationships within sequential data. According to an embodiment, there may be provided a plurality of predetermined transformer blocks, which may be selected by the user.

(18) According to an embodiment, the electronic device **101** may convert the tokens output from the predetermined transformer block from the spatial domain to the frequency domain through the FFT-based frequency domain conversion algorithm. According to an embodiment, the FFT-based frequency domain conversion algorithm may include discrete cosine transform (DCT). In this case, DCT refers to a discrete cosine conversion technique, and since the DCT has the advantage of better energy compaction using a cosine function for frequency decomposition, which is the characteristic of better gathering information into one place, it is widely used in image compression technology. According to an embodiment, the predetermined transformer block may be set by a user. For example, referring to FIG. **3**, assuming that there is a model including 12 transformer blocks, the user may preset first, fourth, seventh, and tenth transformer blocks to perform the FFT-based frequency domain conversion algorithm on the tokens output from the blocks.

(19) In operation **205**, according to various embodiments, the electronic device (e.g., the processor **120** of FIG. **1**) may prune specific tokens based on frequency information of the tokens converted into the frequency domain, and then convert the remaining tokens other than the specific tokens

from the frequency domain to the spatial domain.

(20) According to an embodiment, the electronic device **101** may prune the specific tokens from the tokens based on the frequency information of the tokens converted into the frequency domain. According to an embodiment, the electronic device **101** may prune the specific tokens from the tokens by the number of tokens selected by the user in the order of a token with the highest frequency to a token with the lowest frequency. For example, referring to FIG. 3, the electronic device **101** may prune a token **321** having the highest frequency and a token **322** having the highest frequency selected by the user among tokens **320** converted into a frequency domain output from the first transformer block **311**. For another example, referring to FIG. 3, the electronic device **101** may prune the token **321** having the highest frequency and the token **322** having the second highest frequency by the number of tokens (e.g., two) selected by the user among the tokens **320** output from the predetermined first transformer block **311** and converted into the frequency domain. For another example, referring to FIG. 3, the electronic device **101** may prune the token **331** having the highest frequency and the token **332** having the second highest frequency by the number of tokens (e.g., two) selected by the user among the tokens **330** output from the predetermined second transformer block **312** and converted into the frequency domain. According to an embodiment, the user may set identically or differently the number of tokens to be pruned for each predetermined transformer block.

(21) According to an embodiment, the electronic device **101** may confirm the number of pruned target tokens corresponding to the predetermined transformer block, and may prune the specific tokens from the tokens **320** by number of pruned target tokens from the tokens **320** in the order of a token with the highest frequency to a token with the lowest frequency. The number of pruned target tokens may be preset for each predetermined transformer block. For example, referring to FIG. 3, the number of pruned target tokens corresponding to the predetermined first transformer block **311** may be two, and the number of pruned target tokens corresponding to the predetermined second transformer block **312** may be two. Specifically, referring to FIG. 3, the electronic device **101** may prune the token **321** having the highest frequency and the token **322** having the second highest frequency by the number of pruned target tokens (e.g., two) corresponding to the first transformer block **311** among the tokens **320** output from the predetermined first transformer block **311** and converted to the frequency domain. In addition, the electronic device **101** may prune the token **331** having the highest frequency and the token **332** having the second highest frequency by the number of pruned target tokens (e.g., two) corresponding to the second transformer block **312** among the tokens **320** output from the predetermined second transformer block **312** and converted to the frequency domain.

(22) According to an embodiment, the electronic device **101** may convert the remaining tokens other than the specific tokens from the frequency domain to the spatial domain. For example, referring to FIG. 3, the electronic device **101** may prune the token **321** having the highest frequency and the token **322** having the second highest frequency in the first transformer block **311**, and then convert domains of the seven remaining tokens **323** from the frequency domain to the spatial domain.

(23) According to an embodiment, the electronic device **101** may transmit the remaining tokens converted into the spatial domain to a next transformer block.

(24) According to an embodiment, the electronic device **101** may perform a global average pooling (GAP) layer and an MLP head layer after performing an operation of processing the token using the transformer block.

(25) According to various embodiments, the electronic device that performs the token pruning in the frequency domain may include the processor, and the processor may divide the image frame into the plurality of patches, and may convert the tokens based on the plurality of patches, output from the predetermined transformer block among the plurality of transformer blocks through the fast Fourier transform (FFT)-based frequency domain conversion algorithm from the spatial

domain to the frequency domain, after performing the patch embedding on the plurality of patches, and based on the frequency information of the tokens converted into the frequency domain, after pruning the specific tokens from the tokens, the processor may be configured to convert the remaining tokens other than the specific tokens from the frequency domain to the spatial domain.

(26) According to various embodiments, the processor may be configured to prune the specific tokens from the tokens by the number of tokens selected by the user in the order of a token with the highest frequency to a token with the lowest frequency.

(27) According to various embodiments, the processor may be configured to confirm the number of pruned target tokens corresponding to the predetermined transformer block.

(28) According to various embodiments, the processor may be configured to prune the specific tokens from the tokens by the number of pruned target tokens in the order of a token with the highest frequency to a token with the lowest frequency.

(29) According to various embodiments, a plurality of predetermined transformer blocks may be provided, predetermined transformer blocks may be selected by a user, and as one of the predetermined transformer blocks, the number of pruned target tokens may be set to x in a first transformer block, and as one of the predetermined transformer blocks, the number of pruned target tokens may be set to y in a second transformer block.

(30) According to various embodiments, the processor may be configured to transmit the remaining tokens to the next transformer block.

(31) According to various embodiments, the processor may be configured to perform the patch embedding without generating the CLS token in the process of performing the patch embedding.

(32) According to various embodiments, a method of operating an electronic device for performing token pruning in a frequency domain may include an operation of dividing the image frame into the plurality of patches using the processor of the electronic device, an operation of converting the tokens based on the plurality of patches from the spatial domain to the frequency domain through the fast Fourier transform (FFT)-based frequency domain conversion algorithm, the tokens being output from the predetermined transformer block among the plurality of transformer blocks, after performing patch embedding on the plurality of patches, using the processor, and an operation of converting the remaining tokens other than specific tokens from the frequency domain to the spatial domain, after pruning specific tokens from the tokens based on frequency information of the tokens converted into the frequency domain, using the processor.

(33) According to various embodiments, the operation of pruning the specific tokens may include an operation of pruning the specific tokens from the tokens by the number of the tokens selected by a user in the order of a token with the highest frequency to a token with the lowest frequency.

(34) According to various embodiments, the operation of pruning the specific tokens may include an operation of confirming the number of pruned target tokens corresponding to the predetermined transformer block.

(35) According to various embodiments, the operation of pruning the specific tokens may include an operation of pruning the specific tokens from the tokens by the number of pruned target tokens in the order of a token with the highest frequency to a token with the lowest frequency.

(36) According to various embodiments, the method of operating an electronic device for performing token pruning in a frequency domain may further include an operation of transmitting the remaining tokens to a next transformer block.

(37) According to various embodiments, the operation of performing the patch embedding may include an operation of performing the patch embedding without a CLS token in the process of performing the patch embedding.

(38) The term “module” or “portion” used in the present disclosure may refer to, for example, a unit including one or more combinations of hardware, software, and firmware. For example, the “module” may be interchangeable with a term, such as “logic,” “logical block,” “component,” “circuit,” or the like. The “module” or “portion” may be a minimum unit of a component formed as

one body or a part thereof. The “module” or “portion” may be a minimum unit for performing one or more functions or a part thereof. The “module” or “portion” may be implemented mechanically or electronically. For example, the “module” or “portion” according to an embodiment of the present disclosure may include at least one of an application-specific integrated circuit (ASIC) chip, a field-programmable gate array (FPGA), or a programmable-logic device for performing certain operations which have been known or are to be developed in the future, and may be performed by the processor **120**. According to an embodiment of the present disclosure, at least a part of the device (e.g., modules or their functions) or method (e.g., operations) may be implemented as instructions stored in a computer-readable storage medium (e.g., the memory **130**), e.g., in the form of a program module. The instructions, when executed by a processor (e.g., the processor **120**), may enable the processor to carry out a corresponding function. The computer-readable medium may include, e.g., a hard disk, a floppy disc, a magnetic medium (e.g., magnetic tape), an optical recording medium (e.g., CD-ROM, DVD, magnetic-optical medium (e.g., floptical disk)), or an embedded memory. The instruction may include a code created by a compiler or a code executable by an interpreter. Modules or programming modules in accordance with various embodiments of the present disclosure may include at least one or more of the aforementioned components, omit some of them, or further include other additional components. Operations performed by modules, programming modules or other components in accordance with various embodiments of the present disclosure may be carried out sequentially, in parallel, repeatedly or heuristically, or at least some operations may be executed in a different order or omitted or other operations may be added.

(39) The embodiments disclosed in the present disclosure are presented for explanation and understanding of the disclosed technical content, and are not intended to be limited to the scope of the present disclosure. Therefore, the scope of the present disclosure should be interpreted as including all modifications or various other embodiments based on the technical idea of the present disclosure.

Claims

1. An electronic device for performing token pruning in a frequency domain, the device comprising: a processor, wherein the processor is configured to, divide an image frame into a plurality of patches, convert tokens based on the plurality of patches from a spatial domain to a frequency domain through a fast Fourier transform (FFT)-based frequency domain conversion algorithm, the tokens being output from a predetermined transformer block among a plurality of transformer blocks, after performing patch embedding on the plurality of patches, and convert the remaining tokens other than specific tokens from the frequency domain to the spatial domain, after pruning specific tokens from the tokens based on frequency information of the tokens converted into the frequency domain, wherein the processor is configured to prune the specific tokens from the tokens by the number of the tokens selected by a user in the order of a token with the highest frequency to a token with the lowest frequency.
2. An electronic device for performing token pruning in a frequency domain, the device comprising: a processor, wherein the processor is configured to, divide an image frame into a plurality of patches, convert tokens based on the plurality of patches from a spatial domain to a frequency domain through a fast Fourier transform (FFT)-based frequency domain conversion algorithm, the tokens being output from a predetermined transformer block among a plurality of transformer blocks, after performing patch embedding on the plurality of patches, and convert the remaining tokens other than specific tokens from the frequency domain to the spatial domain, after pruning specific tokens from the tokens based on frequency information of the tokens converted into the frequency domain, wherein the processor is configured to confirm the number of pruned target tokens corresponding to the predetermined transformer block.

3. The electronic device according to claim 2, wherein the processor is configured to prune the specific tokens from the tokens by the number of pruned target tokens in the order of a token with the highest frequency to a token with the lowest frequency.
 4. The electronic device according to claim 3, wherein a plurality of predetermined transformer blocks are provided, the predetermined transformer blocks are selected by a user, as one of the predetermined transformer blocks, the number of pruned target tokens is set to x in a first transformer block, and as one of the predetermined transformer blocks, the number of pruned target tokens is set to y in a second transformer block.
 5. The electronic device according to claim 1, wherein the processor is configured to transmit the remaining tokens to a next transformer block.
 6. The electronic device according to claim 1, wherein the processor is configured to perform the patch embedding without a CLS token in the process of performing the patch embedding.
 7. A method of operating an electronic device for performing token pruning in a frequency domain, the method comprising: dividing an image frame into a plurality of patches using a processor of the electronic device; converting tokens based on the plurality of patches from a spatial domain to a frequency domain through a fast Fourier transform (FFT)-based frequency domain conversion algorithm, the tokens being output from a predetermined transformer block among a plurality of transformer blocks, after performing patch embedding on the plurality of patches, using the processor; and converting the remaining tokens other than specific tokens from the frequency domain to the spatial domain, after pruning specific tokens from the tokens based on frequency information of the tokens converted into the frequency domain, using the processor, wherein the operation of pruning the specific tokens comprises pruning the specific tokens from the tokens by the number of the tokens selected by a user in the order of a token with the highest frequency to a token with the lowest frequency.
 8. The method according to claim 7, wherein the pruning the specific tokens comprises confirming the number of pruned target tokens corresponding to the predetermined transformer block.
 9. The method according to claim 8, wherein the pruning the specific tokens comprises pruning the specific tokens from the tokens by the number of pruned target tokens in the order of a token with the highest frequency to a token with the lowest frequency.
 10. The method according to claim 9, wherein a plurality of predetermined transformer blocks are provided, the predetermined transformer blocks are selected by the user, as one of the predetermined transformer blocks, the number of pruned target tokens is set to x in a first transformer block, and as one of the predetermined transformer blocks, the number of pruned target tokens is set to y in a second transformer block.
 11. The method according to claim 7, further comprising: transmitting the remaining tokens to a next transformer block.
 12. The method according to claim 7, wherein the performing the patch embedding comprises performing the patch embedding without a CLS token in the process of performing the patch embedding.
-