

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12395377
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Wragg; David Paul et al.

Traffic load balancing between a plurality of points of presence of a cloud computing infrastructure

Abstract

Methods and system of traffic load balancing between a plurality of Points of Presence (PoP) of a cloud computing infrastructure are described. A first PoP of multiple PoPs of cloud computing infrastructure that provides a cloud computing service receives a packet. The packet includes as a destination address an anycast address advertised by the first PoP for reaching the cloud computing service. The first PoP identifies a network address of a second PoP that is different from the first PoP. The first PoP forwards the packets as an encapsulated packet to the second PoP to be processed in the second PoP according to the cloud computing service.

Inventors: Wragg; David Paul (London, GB), Guðmundsson; Ólafur (Chevy Chase, MD), Bauer; Lorenz Mathias (London, GB), Fabre; Arthur (London, GB), Majkowski; Marek Przemyslaw (Warsaw, PL)

Applicant: CLOUDFLARE, INC. (San Francisco, CA)

Family ID: 1000008765426

Assignee: CLOUDFLARE, INC. (San Francisco, CA)

Appl. No.: 18/516543

Filed: November 21, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240089315 A1	Mar. 14, 2024

Related U.S. Application Data

continuation parent-doc US 17481181 20210921 US 11824923 child-doc US 18516543
continuation parent-doc US 17023209 20200916 US 11128695 20210921 child-doc US 17481181

Publication Classification

Int. Cl.: H04L67/1008 (20220101); H04L45/00 (20220101); H04L47/122 (20220101); H04L67/1001 (20220101)

U.S. Cl.:

CPC H04L67/1008 (20130101); H04L45/22 (20130101); H04L47/122 (20130101); H04L67/1001 (20220501);

Field of Classification Search

CPC: H04L (47/10); H04L (47/12); H04L (47/122); H04L (47/125); H04L (47/76); H04L (45/22)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7369557	12/2007	Sinha	N/A	N/A
8843998	12/2013	Fu et al.	N/A	N/A
9379981	12/2015	Zhou et al.	N/A	N/A
9467378	12/2015	Stark et al.	N/A	N/A
9584328	12/2016	Graham-Cumming	N/A	N/A
9680951	12/2016	Graham-Cumming	N/A	H04L 45/566
10567333	12/2019	Sawyer et al.	N/A	N/A
10931743	12/2020	Chou et al.	N/A	N/A
11621891	12/2022	Chou et al.	N/A	N/A
2003/0198189	12/2002	Roberts et al.	N/A	N/A
2014/0122698	12/2013	Batrouni	709/224	H04L 41/0896
2015/0215388	12/2014	Kontothanassis et al.	N/A	N/A
2017/0366604	12/2016	McDuff	N/A	N/A
2018/0006952	12/2017	Nakakura et al.	N/A	N/A
2018/0367498	12/2017	Bliss	N/A	H04L 45/74
2019/0208554	12/2018	Ruiz et al.	N/A	N/A
2020/0213154	12/2019	Han et al.	N/A	N/A

OTHER PUBLICATIONS

Eisenbud et al., “Maglev: A Fast and Reliable Software Network Load Balancer”, 13th USENIX Symposium on Networked Systems Design and Implementation (NSDI '16), Available Online at <<https://www.usenix.org/system/files/conference/nsdi16/nsdi16-paper-eisenbud.pdf>>, Mar. 16-18, 2016, pp. 523-535. cited by applicant

Github, “Glb-Director”, Available Online at <<https://github.com/github/glb-director/tree/master/src/glb-redirect>>, Jan. 2020, 1 page. cited by applicant

Herbert et al., “Generic UDP Encapsulation: Draft-ietf-intarea-gue-09”, Internet Draft, Available online at <https://datatracker.ietf.org/doc/draft-ietf-intarea-gue/?include_text=1>, Oct. 26, 2019, 24 pages. cited by applicant

Non-Final Office Action, U.S. Appl. No. 17/481,181, Jan. 6, 2023, 18 pages. cited by applicant

Notice of Allowance, U.S. Appl. No. 17/023,209, May 19, 2021, 9 pages. cited by applicant

Notice of Allowance, U.S. Appl. No. 17/481,181, Jul. 19, 2023, 10 pages. cited by applicant

Olteanu et al., "Stateless Datacenter Load-balancing with Beamer", 15th USENIX Symposium on Networked Systems Design and Implementation, Available Online at <<https://www.usenix.org/system/files/conference/nsdi18/nsdi18-olteanu.pdf>>, Apr. 9-11, 2018, pp. 125-139. cited by applicant

Olteanu et al., "Stateless Datacenter Load-balancing with Beamer", Available Online at <<https://www.usenix.org/conference/nsdi18/presentation/olteanu>>, 2018, 40 slides. cited by applicant

Shirokov et al., "Open-sourcing Katran, a Scalable Network Load Balancer", Open Source, Available Online at <<https://engineering.fb.com/open-source/open-sourcing-katran-a-scalable-network-load-balancer/>>, May 22, 2018, 7 pages. cited by applicant

Unknown author, "What is tunneling or port forwarding", Definition from TechTarget (Year: 2007). cited by applicant

Yong et al., "GRE-in-UDP Encapsulation", RFC: 8086, Internet Engineering Task Force (IETF), Mar. 2017, pp. 1-27. cited by applicant

Primary Examiner: Grijalva Lobos; Boris D

Attorney, Agent or Firm: NICHOLSON DE VOS & ELLIOTT LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a continuation of Ser. No. 17/481,181, filed Sep. 21, 2021, which is a continuation of U.S. application Ser. No. 17/023,209, filed Sep. 16, 2020, now U.S. Pat. No. 11,128,695, which is hereby incorporated by reference.

TECHNICAL FIELD

(1) One or more embodiments relate to the field of network computing; and more specifically, to traffic load balancing between a plurality of Points of Presence of a cloud computing infrastructure.

BACKGROUND ART

(2) Today there is a proliferation of cloud-based services that provide services to users located around the world. To meet a high demand for these services at low latency, a cloud-based service hosts the service on a number of servers located in multiple locations around the world. These locations are referred to as clusters or Points of Presence (PoPs). A PoP includes hardware and software resources for providing a service in a geographical area. A PoP can be part of a data center. The data center can be owned by the service provider or by a third-party providing hardware resources as a service (Infrastructure as a Service) to the service provider.

(3) To ensure the efficient use of the resources within each PoP, a load balancer receives the traffic and distributes the traffic load evenly across the servers of the PoP so that no single server gets overloaded. Thus, a PoP's load balancer steers traffic destined to a given service across a dynamic set of servers. These servers are sometimes referred to as backend servers.

(4) Several solutions for enabling load balancing of traffic within a PoP exist. However, with the increase in that volume of traffic load balancing this traffic within the PoP may not be sufficient as the backend servers may still get overloaded. Additional computing resources may be needed to support this increase in traffic.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The following figures use like reference numbers to refer to like elements. Although the following figures depict various exemplary embodiments, alternative embodiments are within the spirit and scope of the appended claims. In the drawings:
- (2) FIG. 1A illustrates a block diagram of a system including a cloud computing infrastructure in accordance with some embodiments.
- (3) FIG. 1B illustrates a block diagram of an exemplary cloud computing infrastructure for configuring PoP management servers to load balance traffic between PoPs, in accordance with some embodiments.
- (4) FIG. 1C illustrates a block diagram of exemplary inter-PoP forwarding policies that can be configured for a first PoP at a first time, in accordance with some embodiments.
- (5) FIG. 1D illustrates a block diagram of exemplary inter-PoP forwarding policies that can be configured for a first PoP at a second time, in accordance with some embodiments.
- (6) FIG. 1E illustrates a block diagram of an exemplary forwarding unit for forwarding traffic to another PoP according to one embodiment.
- (7) FIG. 1F illustrates a block diagram of an exemplary forwarding unit for forwarding traffic to another PoP according to one embodiment where existing connections are maintained.
- (8) FIG. 2A illustrates a flow diagram of exemplary operations that can be performed for determining forwarding policies for a PoP in the cloud computing infrastructure in accordance with some embodiments.
- (9) FIG. 2B illustrates a flow diagram of exemplary operations that can be performed for load balancing traffic across PoP(s) of a cloud computing infrastructure, in accordance with some embodiments.
- (10) FIG. 2C illustrates a flow diagram of exemplary operations that can be performed for forwarding traffic from one PoP to another while maintaining existing connections.
- (11) FIG. 2D illustrates a flow diagram of exemplary operations that can be performed when receiving a packet from another PoP, in accordance with some embodiments.
- (12) FIG. 3 illustrates a block diagram for an exemplary data processing system that may be used in some embodiments.

DETAILED DESCRIPTION

- (13) The following description describes methods and apparatus for traffic load balancing between a plurality of Points of Presence of a cloud computing infrastructure. In the below description, numerous specific details such as resource partitioning/sharing/duplication embodiments, types and interrelationships of system components, and logic partitioning/integration choices are set forth in order to provide a more thorough understanding. It will be appreciated, however, by one skilled in the art, that the invention may be practiced without such specific details. In other instances, control structures, logic embodiments, opcodes, means to specify operands, and full software instruction sequences have not been shown in detail since those of ordinary skill in the art, with the included descriptions, will be able to implement what is described without undue experimentation.
- (14) References in the specification to “one embodiment,” “an embodiment,” “an example embodiment,” etc., indicate that the embodiment described may include a particular feature, structure, or characteristic, but every embodiment may not necessarily include the particular feature, structure, or characteristic. Moreover, such phrases are not necessarily referring to the same embodiment. Further, when a particular feature, structure, or characteristic is described in connection with an embodiment, it is submitted that it is within the knowledge of one skilled in the art to affect such feature, structure, or characteristic in connection with other embodiments whether or not explicitly described.
- (15) Bracketed text and blocks with dashed borders (e.g., large dashes, small dashes, dot-dash, and dots) may be used herein to illustrate optional operations and/or structures that add additional

features to some embodiments. However, such notation should not be taken to mean that these are the only options or optional operations, and/or that blocks with solid borders are not optional in certain embodiments.

(16) In the following description and claims, the term “coupled,” along with its derivatives, may be used. “Coupled” is used to indicate that two or more elements, which may or may not be in direct physical or electrical contact with each other, co-operate or interact with each other.

(17) Existing Load Balancing Techniques in Cloud Infrastructure

(18) Cloud computing infrastructures can provide services to users through a network computing infrastructure that is distributed across multiple geographical areas. A PoP of a cloud computing infrastructure includes one or more compute servers located at a same geographical area and host the service(s) provided to the users of the cloud computing infrastructure. To address the increase in traffic handled by PoPs of the cloud computing infrastructure and CPU load on these PoPs, the cloud computing infrastructure can increase its capacity by adding hardware (e.g., adding compute servers in existing PoPs, or adding PoPs in a geographical region). However, predicting the future CPU needs of a PoP is challenging, and the lead time for upgrading the hardware and deploying the service in a PoP is long. A solution to addressing the increase of network traffic is to regulate the CPU load across the PoPs of the cloud computing infrastructure. One mechanism of doing so is the load balancing of the CPU load between the existing PoPs of the cloud computing infrastructures. Some solutions exist for load balancing by managing the network traffic arriving at a PoP. One such solution relies on a routing and reachability protocol, e.g., Border Gateway Protocol (BGP). For example, in an anycast routing based cloud computing infrastructure, where the traffic that reaches a compute server of a PoP is determined by the anycast Internet Protocol addresses that the compute server advertises to the network through BGP, one solution of load balancing across PoPs is based on withdrawing some of the advertised IP addresses from compute servers of a PoP. The withdrawal of these IP addresses reduces the traffic reaching the compute servers of the PoP (as these servers will no longer receive traffic destined to the withdrawn IP addresses), and consequently reduces the compute processing (CPU) load on the PoP. However, this BGP-based approach has several drawbacks. The solution works at the granularity of the anycast IP addresses that are advertised from the PoPs. Therefore, changes to the cloud computing infrastructure are not fine-grained, and there are limits to the ability of the cloud computing infrastructure to predict the resulting impact of these changes on CPU load. Further, when a PoP withdraws or restores the advertisement for an IP address, long lived Transmission Control Protocol (TCP) connections (or User Datagram Protocol (UDP)-based pseudo-connections) destined for that IP address will be broken. Another consequence of withdrawal of the IP addresses is that traffic associated with a withdrawn IP address will go to some other PoPs instead. However, there is no control over which PoPs are to receive the traffic. The new PoP to receive the traffic destined to the IP address is determined based on the routing protocol (e.g., the anycast protocol). In addition, changes to the routing network in response to BGP advertisement take some time (e.g., a minute or two) to propagate across the network and to take full effect. Therefore, there is a need for a more efficient load-balancing mechanism between PoPs of a cloud computing infrastructure that provides more control over the re-routing of the traffic.

(19) In one embodiment, a first PoP of a plurality of PoPs that form a cloud computing infrastructure receives a packet. The packet includes as a destination address an anycast address advertised by the first PoP for reaching a cloud computing service. In some embodiments, the first PoP may optionally determine whether the packet is a security threat. Upon determining that the packet is a security threat, the packet is dropped. Alternatively, upon determining that the packet is not a security threat, the first PoP identifies, based on a forwarding table entry and an identifier for a flow of the packet, a compute server that is part of a second PoP of the plurality of PoPs that is different from the first PoP. Upon identification of the compute server, the first PoP encapsulates the packets into an encapsulated packet that includes an external address of the compute server as

the destination address of the packet. The first PoP submits the encapsulated packet to the second compute server to be processed in the second PoP according to the cloud computing service. In this embodiment, while the packet is received at a first PoP based on an anycast address and the first PoP is the destination of the anycast address, the packet is forwarded to be processed to another PoP that is remote and different from the first PoP according to the forwarding table entry. This mechanism enables load balancing of traffic across multiple PoPs.

(20) In another embodiment, the load balancing of traffic across multiple PoPs is further performed by avoiding disruption of existing connection(s) in the PoPs. In one embodiment, a first PoP receives a first packet that includes as a first destination address an anycast address advertised by the first PoP for reaching the cloud computing service. The first PoP identifies, based on a forwarding table entry and a first identifier for a first flow of the first packet, a first compute server that is local to the first PoP and a second compute server that is part of a second PoP of the plurality of PoPs that is different from the first PoP. In response to determining that a first connection is established at the first compute server for the first flow, the first packet is processed at the first compute server according to the cloud computing service. Alternatively, in response to determining that no connection is established at the first compute server for the first flow, the first PoP determines whether the first packet is an initial packet of a new connection for the first flow. When it is determined that the first packet is not the initial packet of the new connection, the first packet is transmitted to the second compute server that is part of the second PoP of the plurality of PoPs to be processed in the second PoP according to the cloud computing service.

(21) In some embodiments, when it is determined that the first packet is the initial packet of a new connection, a determination of a forwarding direction is performed. When it is determined that the forwarding direction indicates a local forwarding, the first packet is processed at the first compute server of the first PoP; and when it is determined that the forwarding direction indicates a remote forwarding, the first packet is processed at the second compute server of the second PoP.

(22) The embodiments described herein provide several advantages when compared to existing load balancing solutions. For example, in contrast to other solutions, such as BGP-based load balancing solutions, the embodiments herein maintain lingering connections. That is, changes made to the forwarding policy of a PoP will not break established connections while enabling new connections to obey the new forwarding policy. Further, the embodiments herein enable a more fine-grained and explicit control of forwarding policies from a PoP to another PoP. Thus, these embodiments offer predictable control of the destination of traffic and consequently more predictable control of resource usage in the cloud computing infrastructure. In some embodiments, encapsulation protocol(s) can be used to forward packets between PoPs. The use of encapsulation tunnels to redirect traffic from a PoP to another PoP enables a fine-grained control of the traffic forwarding. A PoP may forward parts of its incoming traffic to one or multiple other PoPs. For example, a forwarding policy can specify the percentage of traffic to be forwarded to each target PoP for a given type of traffic. The combination of fine-grained control and unobtrusive changes to forwarding policies allow the systems herein to make frequent changes to forwarding policies and forwarding table entries, which prevent PoPs from becoming overloaded.

(23) FIG. 1A illustrates a block diagram of a system **100** including a cloud computing infrastructure, in accordance with some embodiments. The cloud computing infrastructure can provide different services to users. The cloud computing infrastructure can provide services such as protecting against internet-based threats, providing performance services for customers (e.g., acting as a content delivery network (CDN) and dynamically caching customer's files close to visitors, page acceleration, etc.), TCP stack optimizations, and/or other services. The system **100** includes the client devices **110A-L**, the compute servers **120A-N** included in different Points of Presence (PoPs) **150A-M**, the server **112**, and the optional origin servers **130A-K**.

(24) Each one of the client devices **110A-N** is a computing device (e.g., laptop, workstation, smartphone, mobile phone, tablet, gaming system, set top box, wearable device, Internet of Things

(IoT) device, etc.) that is capable of transmitting and/or receiving network traffic. Each client device may execute a client network application such as a web browser, native application, or other application that can access network resources (e.g., web pages, images, word processing documents, PDF files, movie files, music files, or other computer files). A client device, e.g., client device **110A**, is operative to transmit requests for network resources through the network **105**. The requests are transmitted in one or more packets (e.g., IP/TCP packets) **101A** to a compute server, e.g., compute server **120A**, of the cloud computing infrastructure **101**. The client device **110A** is operative to receive response(s) from a compute server of the cloud computing infrastructure **101**. In some embodiments, the responses are received from another compute server, e.g., compute server **120B**, that is different from the one to which the requests are transmitted. The embodiments described herein will use the client device **110A** as an exemplary client device, however, one of ordinary skill in the art would understand that multiple client devices can transmit requests and receive responses to/from the PoPs of the cloud computing infrastructure **101**. Typically, multiple client devices communicate with each one of the compute servers of a PoP of the cloud computing infrastructure **101**.

(25) The compute servers **120A-N** are part of the cloud computing infrastructure. The compute servers **120A-N** are physical electronic devices and can be geographically distributed (e.g., in different locations throughout the world). There may be hundreds of compute servers as part of the cloud computing infrastructure **101**. The compute servers **120A-N** may be part of PoPs **150A-M**. Each one of the PoPs may include one or more of the compute servers. For example, PoP **150A** includes the compute servers **120A-P**, PoP **150B** includes the compute servers **120B-Q**, and PoPs **150R-M** includes the compute servers **120R-N**. Each one of the compute servers **120A-N** may host a service that is offered by the cloud computing infrastructure **101** to customers or users. The compute servers **120A-N** may also host services and/or functionalities that are used to support the service offered by the same or other compute servers in a PoP. For example, a compute server may implement control functionalities, DNS service (e.g., one or more authoritative name servers, one or more proxy DNS servers), a proxy service, a CDN service, TCP stack optimization, a caching service, etc. Each PoP may also include one or more other pieces of network equipment such as router(s), switch(es), hub(s), and/or other intermediate computing devices. Each PoP (and each compute server) may be part of a different data center and/or colocation site.

(26) In some embodiments, the compute server **120A** is operative to host the service(s) **123A** of the cloud computing infrastructure **101**. In these embodiments, upon receipt of a request for a service **123A**, the compute server **120A** is operative to process the received request at the network layer (Layers **3/4**) and further operative to fulfill the request for service at the application layer (Layer **7**).

(27) In some embodiments, the PoP **150A** includes a forwarding unit **124A**. The forwarding unit **124A** is operative to receive packets from one or more client devices and process the packets according to multiple embodiments described herein. In one embodiment, the forwarding unit **124A** is operative to transmit the packet to a local compute server, e.g., compute server **120A** included in PoP **150A**, to be processed at the first PoP **150A**. In this embodiment, a request for a service hosted by the cloud computing infrastructure is fulfilled at the first PoP **150A**. In another embodiment, the forwarding unit **124A** is operative to transmit the packet to a remote compute server of a second PoP that is different from the first PoP. For example, the forwarding unit **124A** is operative to transmit the packet to the compute server **120B** of PoP **150B**. In some embodiments, the forwarding unit **124A** may be operative to transmit all packets to a local compute server of the first PoP **150A** and based on certain criteria the packet may be processed at the local compute server or alternatively forwarded to the second compute server of a second PoP **150B**. In some embodiments, the forwarding unit **124A** is implemented in a compute server of a PoP and is part of a local network of the PoP. In some embodiments, the forwarding unit **124A** is operative to receive traffic that is destined to an anycast address (e.g., anycast Internet Protocol (IP) address) that is advertised by the PoP **150A**. For example, the anycast address advertised by the PoP can be

advertised through a BGP protocol implemented in a router that is part of the PoP **150A**. In some embodiments, the forwarding unit **124A** and the router can be implemented on the same network device. Alternatively, the forwarding unit **124A** and the router can be implemented on separate network devices that are part of the same local network in the PoP **150A**. In some embodiments, the anycast address can be associated with a network resource (e.g., a domain) that is hosted by an origin server.

(28) The compute servers **120A-N** can be geographically distributed across the multiple PoPs, which decreases the distance between requesting client devices and the compute servers and decreases the time necessary to respond to a request. In some embodiments, the compute servers **120A-N** may operate as reverse proxies and may receive requests for network resources (e.g., HTTP requests) of a domain of an origin server **130A** included in packets. The particular compute server that receives a request from a client device may be determined by the network infrastructure according to an anycast implementation or by a geographical load balancer. For instance, the compute server **120A**, the compute server **120P**, and the computer server **120N** may have a same anycast address for a domain of the origin server **130A**. If the origin server **130A** handles the domain “example.com,” a domain name service (DNS) request for “example.com” returns an address record having the anycast IP address of the compute server **120A**, compute server **120P**, and computer server **120N**. Which one of the compute server **120A**, compute server **120P**, and compute server **120N** receives a request from a client device depends on which compute server is closest to the client device in terms of routing protocol configuration (e.g., BGP configuration) according to an anycast implementation as determined by the network infrastructure (e.g., router(s), switch(es), and/or other network equipment between the requesting client and the compute servers **120A**, **120P**, and **120N**). By way of example, the client device **110A** is closest to the compute server **120A**. Accordingly, requests from the client device **110A** are received at the compute server **120A**. While at a moment in time traffic from a client device transmitted to the cloud computing infrastructure is received at the first PoP for the anycast address, at another moment in time, the traffic from the same client device and for the same anycast address can be received at another PoP due to the anycast protocol. Further, traffic destined to the same anycast address that originates from different client devices may be routed to the same PoP of the cloud computing infrastructure or to different PoPs of the cloud computing infrastructure. The number of client devices and compute servers illustrated in FIG. 1A is exemplary. The distributed cloud computing network may include hundreds to thousands (or more) compute servers and each compute server may receive requests from thousands or more client devices.

(29) In some embodiments, a compute server in a PoP is operative to send heartbeat packets to signal node liveness to the PoP management server **140A**. The compute server is further operative to perform service health checks on the node on which it is implemented. The compute server is operative to process packets at the network layer and at the application layer for providing the cloud computing service.

(30) Each one of the compute servers has an external IP address (sometimes referred to as a global IP address). In some embodiments, each one of the compute servers has an anycast IP address in addition to the external IP address. The anycast IP address is assigned to multiple compute servers located in different geographical regions and which are part of multiple PoPs of the cloud computing infrastructure. For example, the compute server **120A** has a first external IP address that is unique to the computer server **120A** and a first anycast IP address; the compute server **120B** has a second external IP address that is unique to the compute server **120B** and a second anycast IP address. In some embodiments, the compute server **120A** and the compute server **120B** have the same anycast IP address. In other embodiments, the compute server **120A** and the compute server **120B** have different anycast IP addresses.

(31) Each one of the PoPs includes one or more PoP management servers **140A**. A PoP management server is operative to track the state of the compute servers in the PoP, receive and

store forwarding policies, manage and update forwarding tables of the compute servers based on the forwarding policies, manage a prefix list, and/or perform other networking related control operations. The PoP management server, e.g., PoP management server **140A**, is operative to receive and respond to control commands from a central controller **114** of the cloud computing infrastructure **101**. In some embodiments, the PoP management server **140A** is operative to receive and respond to control commands that cause the configuration, update, and query of inter-PoP forwarding policies. As it will be described in further detail below, the forwarding policies enable the dynamic load balancing of traffic across multiple PoPs of the system **100**.

(32) The server **112** is operated by the cloud computing infrastructure **101** and provides a set of tools and interfaces for a customer to, among other things, configure code to run in the PoPs **150A-M** of the cloud computing infrastructure **101**. The server **112** includes a central controller **114**. The central controller **114** is operative to configure, query, and update forwarding policies in the PoPs. In some embodiments, the central controller **114** has a global view of the cloud computing infrastructure that includes the multiple PoPs and is operative to configure and/or update forwarding policies within the PoP that cause traffic received at a first PoP to be forwarded to another PoP to be processed.

(33) In an embodiment, the server **112** allows a customer to provision one or more services (e.g., malicious domain emulation service, API compatibility service, a proxy service, a CDN service, TCP stack optimization, a caching service, etc.) to the cloud computing infrastructure through DNS. For example, DNS record(s) of a customer are changed such that DNS records of hostnames point to an IP address of a compute server instead of the origin server. In some embodiments, the authoritative name server of the customer's domain is changed to an authoritative name server of the service and/or individual DNS records are changed to point to the compute server (or point to other domain(s) that point to a compute server of the service). For example, the customers may change their DNS records to point to a CNAME record that points to a compute server of the service. In one embodiment, customers may use the server **112** to change their authoritative name server to an authoritative name server of the cloud computing infrastructure and/or change their zone file records to have their domain point to the compute servers. Pointing the DNS or similar records to the compute server **120A** enables the compute server to act as a proxy for requests directed to the origin server. In this manner, the requests for content items sent to the origin server and the responses thereto pass through the computer server and enable the one or more worker scripts to be executed. While some embodiments are described with the server **112** performing the provisioning of the services in addition to the control and configuration of the forwarding policies, in other embodiments, the provisioning of the services can be performed by another server that is separate from the server **112**. While the embodiments herein will be described with respect to a proxy service provided by the cloud computing infrastructure **101**, where a compute server of the cloud computing infrastructure is to receive a request for a network resource instead of an origin server that hosts the network resource, in other embodiments, the type of services that are offered by the cloud computing infrastructure can vary without departing from the scope of the embodiments described herein.

(34) The origin servers **130A-K**, which may be owned or operated directly or indirectly by customers of the cloud computing infrastructure **101**, are computing devices on which a network resource resides, is generated, and/or otherwise originates (e.g., web pages, images, word processing documents, PDF files movie files, music files, or other content items). Although FIG. **1A** illustrates the origin server **130A** communicating with the compute server **120A**, the origin server **130A** may also communicate with one or more of the other compute servers **120**. Similarly, while FIG. **1A** illustrates a case with the PoP **150A** servicing a single origin server **130A**, the PoP **150A** can service any number of origin servers.

(35) The cloud computing infrastructure **101** enables load balancing of traffic across the multiple PoPs. In some embodiments, the receipt of a flow of traffic in a PoP of the cloud computing

infrastructure **101** is performed based on an anycast protocol. The system of FIG. 1A enables the forwarding of requests from a PoP that receives the request according to the anycast protocol (which generally is the PoP closest to the client device) to another PoP of the cloud computing infrastructure **101**. The forwarding of traffic to the other PoP can be due to the fact that the CPU usage of the first PoP approaches a load limit.

(36) In operation, first packets **101A** including a first request for a first resource are transmitted from client device **110A** to the PoP **150A**. For example, the first packets **101A** may be IP packets destined to a first anycast address advertised by the PoP **150A**. In some embodiments, the first anycast address can be advertised by a router of the PoP **150A** through a BGP protocol to one or more other routers in the network **105**. The packets **101A** may include a request (e.g., HTTP GET request, etc.) for a network resource. The packets **101A** are received at the PoP **150A** as a result of an anycast protocol selecting the PoP **150A** from a set of PoPs that have advertised the first anycast IP address in the network **105**. The PoP **150A** determines based on a forwarding table entry and an identifier of the flow of the packets that the packets are to be processed locally. In response to determining that the packets **101A** are to be processed locally, the packets are processed at a compute server **120A**. The compute server **120A** fulfills the request included in the packets **101A**. In some embodiments, the compute server **120A** may transmit the packets **103A** to an origin server **130A** for fulfilling the request. In these embodiments, the compute server **120A** receives the response included in the packets **104A**. The response includes the requested network resource. Alternatively, the compute server **120B** may retrieve the network resource from a cache (not illustrated) without transmitting the request to the origin server **130A**. The compute server **120A** then transmits the network resource to the client device **110A** in a set of one or more packets **105A**.

(37) In another example, when it is determined that one or more compute servers of the PoP **150A** are overloaded or any other criteria is met, as it will be discussed in further detail below, the PoP **150A** can be configured to forward traffic to another PoP, e.g., PoP **150B**. Thus, when packets **101B** including another request for the service (e.g., a request for a network resource such as an HTTP GET), are received at the PoP **150A**, the packets **101B** are transmitted to the PoP **150B** to be processed instead of being processed locally. The second PoP **150B** is different from the PoP **150A**. The second PoP **150B** receives the packets. The second compute server **120B** processes the packets at the network layer and at the application layer causing the request for the network resource to be fulfilled at the second compute server **120B** instead of the first compute server which originally received the request from the client device **110A**. For example, the compute server **120B** may transmit the request to the origin server **130A** and receive a response including the network resource from the origin server. Alternatively, the compute server **120B** may retrieve the network resource from a cache (not illustrated) without transmitting the request to the origin server **130A**. The compute server **120A** then transmits the network resource to the client device **110A** in a set of one or more packets **105B**. In some embodiments, the second compute server **120B** may forward the packets to another PoP (e.g., a third PoP that is different from the first PoP **150A**) instead of processing the packet at the application layer. In these embodiments, the packets are processed at the application layer at the other PoP. In some embodiments, the packets may be successively forwarded to multiple PoPs prior to reaching a target PoP that is to process the packets at the application layer.

(38) In some embodiments, once the request is processed at the compute server **120B**, the compute server transmits the response to the client device **110A**. In some embodiments, the response is transmitted through a direct server return to the client device, i.e., the response does not need to be forwarded back to the first compute server **120A** to return to the client device. In other embodiments, the response is transmitted to the client device **110A** via the first compute server **120A**. In these embodiments, the response is transmitted to the first compute server **120A** and the first compute server **120A** forwards the response to the client device **110A**. In some embodiments, the response is encapsulated to be transmitted through the tunnel between the second compute

server **120B** and the first compute server **120A**. The first compute server decapsulates the packets including the response and transmits the packets to the client device.

(39) In some embodiments described herein, the path of outgoing traffic (from the PoP to the client devices) is unchanged when compared with the path of outgoing traffic in existing cloud computing infrastructures. The outgoing traffic goes from a compute server to the client devices through the network **105** (e.g., the Internet). In contrast, the incoming traffic (from client device to PoP) may be forwarded to one or more other PoP prior to being processed. In some embodiments, incoming traffic may represent a minority of the traffic received at the PoP and the additional load on the cloud infrastructure of forwarding traffic to another PoP does not overload the system or create a problem. The extra hop (from PoP **150A** to PoP **150B**) does not cause a significant increase in latency for client devices accessing the cloud computing service, or an undue increase in the CPU load on the compute servers of the cloud computing infrastructure.

(40) Determination of Forwarding Policies

(41) FIG. **1B** illustrates a block diagram of an exemplary cloud computing infrastructure for configuring PoP management servers to load balance traffic between PoPs, in accordance with some embodiments. The server **112** includes the central controller **114** and an analytics unit **116**.

(42) The analytics unit **116** is operative to 1) receive feedback from the compute servers in PoPs of the cloud computing infrastructure **101**; 2) identify one or more PoPs from which traffic needs to be forwarded to another PoP, and 3) identify one or more PoPs that can receive traffic. In some embodiments, the feedback is periodically received from the compute servers of the PoPs. For example, each one of the compute servers from the PoPs **150A-N** may be configured to push to the server **112** the feedback at a regular time interval. In an alternative embodiment, the server **112** may be configured to request feedback from the compute servers of the PoPs **150A-N**. In some embodiments, the feedback of a compute server can be either transmitted from the compute server to the server **112** or alternatively, the feedback may be gathered at the PoP management server, e.g., at the PoP management server **140A**, and transmitted from the PoP management server to the server **112**. The feedback of a compute server is indicative of a state of the compute server at a given time. In some embodiments, the feedback can be a parameter indicative of a CPU load of the compute server. Alternatively or additionally, the feedback can be a parameter that measures the Quality of Service (QoS) provided by a service implemented on the compute server. The feedback can be a measure of latency and/or reliability of the service at the compute server.

(43) The analytics unit **116** is operative to identify one or more PoPs for which traffic needs to be forwarded to another PoP. In some embodiments, the identification of a PoP can be performed based on the feedback received for one or more of the compute servers of the PoP. For example, the identification of the PoP can be performed based on the CPU load of the compute servers of the PoP. In some embodiments, PoP **150A** can be identified as a PoP from which traffic needs to be forwarded to another PoP when at least one of the compute server's CPU load exceeds a threshold value. In other embodiments, PoP **150A** can be identified when the CPU load of each one of multiple compute servers (e.g., a number *N* of compute servers, a percentage of the compute servers, etc.) is greater than the threshold value. In other embodiments, PoP **150A** can be identified when a combined CPU load of multiple compute servers of the PoP **150A** exceeds a threshold value. Alternatively or additionally, the identification of the PoP can be performed based on other criteria such as QoS, or type of service. For example, traffic of a particular type of service (e.g., video service) may have stricter requirements for latency/reliability compared to other types of services (e.g., regular web page surfing, email, etc.). Thus, a PoP that is receiving the type of traffic with stricter requirements may not be identified as a PoP from which traffic is to be forwarded. In other embodiments, the PoP **150A** may be selected and the type of service may be used in the determination of forwarding policies and/or forwarding table entries for the particular PoP. In another example, the identification of the PoP can be performed based on other parameters than qualify the service provided at the PoPs (e.g., Quality of Service (QoS) parameters, latency and/or

reliability measures, etc.).

(44) The analytics unit **116** is further operative to identify one or more PoPs to which traffic can be forwarded from one or more other PoPs. In some embodiments, a PoP can be identified, e.g., PoP **150B**, based on the feedback received from one or more of the compute servers of the PoP. For example, the PoP **150B** is identified based on the feedback that is received for the compute servers **120B-Q** of the PoP **150B**. The feedback received from the compute servers can be a parameter indicative of the CPU load. In some embodiments, the analytics unit **116** considers all of the compute servers **120B-Q** of the PoP **150B**. For example, the analytics unit **116** may determine that individual compute servers have a CPU load that is lower than a predetermined threshold value. In some embodiments, the analytics unit **116** may evaluate whether at least one of the compute servers in the PoP **150B** has a CPU load that is lower than the threshold value to identify the PoP **150B** as a PoP that can receive forwarded traffic. Alternatively, the analytics unit **116** may identify the PoP **150B** when at least a number of compute servers (a predetermined number of compute servers, or a percentage of the compute servers) have a CPU load that is lower than the threshold value. Alternatively or additionally, the analytics may determine that a combination of the compute servers **120B-Q** (e.g., a subset or all of the compute servers **120B-Q**) have a combined CPU load that is lower than a predetermined threshold value. In these embodiments, the CPU load of some of the compute servers of the PoP **150B** may have a relatively high load but a total load across the compute servers in the PoP does not exceed a threshold value. Alternatively or additionally, the identification of the PoP can be performed based on other criteria such as QoS or type of service. For example, traffic of a particular type of service (e.g., video service) may have stricter requirements for latency/reliability compared to other types of services (e.g., regular web page surfing, email, etc.). Thus, a PoP that is receiving the type of traffic with stricter requirements may not be identified as a PoP to which traffic is to be forwarded. While the example of FIG. **1B** shows a PoP **150A** and PoP **150B** being identified as respectively a PoP from which traffic is to be forwarded and a PoP to which traffic is to be forwarded, in other embodiments, multiple PoPs can be identified as either PoPs from which traffic is to be forwarded or as PoPs to which traffic can be forwarded.

(45) Once the PoP(s) from which traffic is to be forwarded and the PoP(s) to which traffic is to be forwarded are identified, the central controller **114** determines the inter-PoP forwarding policies **142**. The inter-PoP forwarding policies **142** determine the forwarding of traffic between the PoPs that enables the load balancing of traffic between the PoPs of the cloud computing infrastructure **101**. A forwarding policy for a PoP, e.g., PoP **150A**, includes 1) an amount of traffic that needs to be forwarded from the PoP and 2) a destination network address of a PoP towards which traffic is to be forwarded. In some embodiments, the destination network address can be an external address (e.g., external IP address) of a compute server of the PoP, where the compute server hosts the cloud computing service. Alternatively or additionally, the destination network address can be an external address of a network device in the PoP that has other functionalities. For example, the network device can be a router or a forwarding unit in the PoP. In some embodiments, the destination address is a virtual IP address that is routed to any one of a group of one or more compute servers of the PoP.

(46) In some embodiments, the determination of the forwarding policies is performed based on the identified one or more PoP(s) from which traffic needs to be forwarded and based on the identified one or more PoP(s) towards which traffic needs to be forwarded. In some embodiments, the determination of the forwarding policies is performed by selecting a pair of PoPs, a first PoP **150A** and a second PoP **150B**, where the first PoP is the PoP from which traffic is to be forwarded and the second PoP **150B** is the PoP that is to receive the traffic. In some embodiments, the determination of the forwarding policies for PoP **150A** includes the selection of more than one PoP to which traffic is to be forwarded from PoP **150A**. In some embodiments, the forwarding policies determined enable the PoP **150A** to forward traffic to PoP **150B** and further receive traffic from

another PoP. In some embodiments, the PoPs can be located in a same geographical region, e.g., first region **102A**, than the PoP of the compute server **120A**, while in other embodiments, the PoP(s) can be located in different geographical regions, e.g., first region **102A** and second region **102B**.

(47) The server **112** transmits the forwarding policies **142** to one or more PoPs of the cloud computing infrastructure **101**. For example, the server **112** transmits the forwarding policies to PoP **150A** for forwarding traffic to PoP **150B**. The configuration of the PoPs results in at least one PoP, e.g., PoP **150A**, forwarding a portion of the traffic received according to the anycast protocol to another PoP, e.g., **150B**. In some embodiments, the forwarding policies for a PoP may cause the PoP to forward traffic to multiple ones of the other PoPs in the cloud computing infrastructure **101**.

(48) In some embodiments, PoP **150A** may be configured to forward traffic to PoP **150B** and PoP **150B** may further be configured to forward traffic to another PoP, such as PoP **150R**. This configuration can be used to avoid undesired latency impact for users. For example, if PoP **150A** is identified as an overloaded PoP that needs to forward traffic to another PoP and PoP **150R** is identified as an underloaded PoP, traffic from PoP **150A** could be sent to PoP **150R**. However, PoP **150A** and **150R** are located in two different regions, region **102A** and region **102B**, and PoP **150A** is determined to be far from PoP **150R**. In contrast, PoP **150B** can be closer to PoP **150A** and to PoP **150R** as a distance between PoP **150A** and PoP **150B** and the distance between PoP **150B** and PoP **150R** are less than the distance between PoP **150A** and **150B**. In this scenario, the controller **114** determines one or more forwarding policies that cause traffic to be forwarded from PoP **150A** towards PoP **150B** and traffic to be forwarded from PoP **150B** to PoP **150R**. This ensures that none of PoP **150A** and PoP **150B** get overloaded while avoiding a latency overhead for packets if forwarded from PoP **150A** to PoP **150P**. The determined forwarding policies reduce the maximum latency impact encountered by users.

(49) The determination of the inter-PoP policies **142** is a dynamic mechanism that is performed periodically based on the feedback received for the compute servers of the PoPs. In some embodiments, changes to the forwarding policies can occur at a high frequency such as multiple times per hour. In some embodiments, the inter-PoP forwarding policy that is determined, can be a new inter-PoP forwarding policy (e.g., a policy to forward traffic from PoP **150A** to PoP **150B**), an update to an existing inter-PoP policy, or a cancellation of a previously existing inter-PoP forwarding policy. For example, when the central controller **114** determines that an existing inter-PoP forwarding policy is to be updated, it may transmit a request to modify the fraction of traffic that is forwarded for an established forwarding relationship.

(50) FIG. **1C** illustrates a block diagram of exemplary inter-PoP forwarding policies **142A** that can be configured for PoP **150A** at a first time **t1**, in accordance with some embodiments. FIG. **1C** shows two scenarios of forwarding policies that can be configured for a PoP of the cloud computing infrastructure **101**. In scenario A, the forwarding policy **146A** identifies an amount of traffic that is to be forwarded for the PoP **105A**. The amount of traffic can be expressed in percentage of traffic received from client devices by the PoP **150A**. The amount of traffic is associated with a destination address, e.g., address of PoP **150B**. For example, the destination address can be an external IP address of a compute server **120B** in PoP **150B**. Optionally, inter-PoP forwarding policies **142A** can include an additional forwarding policy **147A** that specifies an additional amount of traffic (e.g., 10%) that is to be forwarded to a destination address of another PoP, e.g., PoP **150C** (not illustrated in FIG. **1A**). When multiple forwarding policies are defined, this results in a total amount of traffic to be forwarded to multiple PoPs from PoP **150A**, here 15% of traffic is forwarded, while the remainder of the traffic (e.g., 85% of traffic received at PoP **150A**) is not forwarded and is processed locally by one or more compute servers of the PoP **150A**. In scenario A, one or both of the forwarding policies **146A** and **147A** can be updated at time **t2**. FIG. **1D** illustrates the updated forwarding policies at the second time **t2**. In this example, the forwarding policy **146A** is updated to increase the amount of traffic that is forwarded to PoP **150B** from 5% to

10%, while the amount of traffic that is forwarded to PoP **150C** is not increased and remains the same. The increase of the amount of traffic that is to be forwarded to PoP **150B** results in a total amount of traffic to be forwarded to multiple PoPs from PoP **150A** to increase, here 20% of traffic is forwarded. The remainder of the traffic (e.g., 80% of traffic received at PoP **150A**) is not forwarded and is processed locally by one or more compute servers of the PoP **150A**.

(51) Returning to FIG. **1C**, in another scenario, scenario B, the forwarding policies **142A** may include forwarding policies, e.g., **148A** and **149A**, that specify an amount of traffic to be forwarded for each type of traffic from multiple types of traffic. For example, the traffic can be classified in different tiers, where a first tier (Tier one) is associated with a first type of service provided at the cloud computing infrastructure **101** and a second tier (Tier two) that is associated with a second type of service that is provided at the cloud computing infrastructure **101**. In one non-limiting example, a first tier of traffic can be traffic of customers that subscribe to a paying service while the second tier can be traffic of customers that subscribe to a free service. Additional tiers of traffic can be defined (e.g., each tier can be associated with a different level of paid service(s), etc.) resulting in an amount of traffic for each tier being forwarded from PoP **150A** to another PoP. In some embodiments, some tiers of traffic may not be forwarded (e.g., traffic of customers that subscribe to paid service) while other tiers of traffic are forwarded (e.g., traffic of customers that subscribe to free service). In the illustrated example of FIG. **1C**, policy **148A** defines that 2% of tier one of traffic is to be forwarded to PoP **150B** and 5% of tier two of traffic is to be forwarded to PoP **150B** at time **t1**, while the remainder of the traffic (e.g., 93% of traffic received at PoP **150A**, which may include tier one or tier two traffic) is not forwarded and is processed locally by one or more compute servers of the PoP **150A**. The forwarding policies are later updated (at time **t2**). FIG. **1D** illustrates the updated forwarding policies. The forwarding policies are updated such that 0% of tier one traffic is to be forwarded, indicating that the traffic in this first tier is not forwarded and is processed locally in the PoP **150A**. Forwarding policy **149A** is updated at time **t2** such that 10% of tier two traffic is forwarded to PoP **150B**. The decrease in the amount of tier one traffic that is to be forwarded from PoP **150A** and the increase in the amount of tier two traffic that is to be forwarded from PoP **150A** results in a total amount of traffic to be forwarded from PoP **150A** to increase, here from 7% to 10% of traffic is forwarded. The remainder of the traffic (e.g., 90% of traffic received at PoP **150A**) is not forwarded and is processed locally by one or more compute servers of the PoP **150A**. The embodiments described herein allow for a fine-grained control over the amount of traffic being forwarded per tier of traffic handled by a PoP.

(52) Configuration of Forwarding Tables Based on Forwarding Policies.

(53) Referring back to FIG. **1B**, once the inter-PoP forwarding policies **142** are transmitted from the server **112** to the PoP(s) **150A-N**, they are converted into forwarding table entries in the PoP(s). For example, a PoP management server **140A** receives and stores the inter-PoP policies **142A**. The inter-PoP forwarding policies **142A** are converted into forwarding table entries **143A**. The PoP controller **140A** updates the forwarding tables in response to changes to the inter-PoP forwarding policies. In some embodiments, in order to minimize disruption to established connections, the PoP controller **140A** modifies as many forwarding table entries as are needed to bring the forwarding table into conformance with the forwarding policy. The PoP controller **140A** selects from table entries in the forwarding table(s) **126A**, one or more table entries to be modified. When selecting forwarding table entries to modify, the PoP controller **140A** may select those that were least recently modified, which allows to minimize disruption of established connections.

(54) In some embodiments, the forwarding table entries **143A** are transmitted to one or more of compute servers **120A-P** of the PoP **150A** to be stored in forwarding tables of the compute servers. In some embodiments, different forwarding entries can be generated for different compute servers of the same PoP. For example, if the anycast address associated with a first compute server **120A** is different from the anycast address associated with a second compute server **120B**, different traffic will reach the different compute servers based on these two different anycast addresses. In this

embodiment, different forwarding entries are populated for the different compute servers. In other embodiments, the same forwarding table entries are forwarded to multiple compute servers in a PoP.

(55) In some embodiments, the forwarding table entries **143A** are stored in the forwarding unit **124A**. In these embodiments, the packets received at the PoP are processed at the forwarding unit **124A** prior to being transmitted according to the forwarding table entries to a local compute server of the PoP or a remote compute server of another PoP to be processed according to the cloud computing service.

(56) While a single forwarding table **126A** is illustrated in the forwarding unit **124A**, in some embodiments multiple forwarding tables can be configured. For example, a first forwarding table can be configured for a first tier of traffic (tier one) and a second forwarding table can be configured for another tier of traffic (e.g., tier two). In some exemplary embodiments, a forwarding table is configured for each tier of traffic. In other exemplary embodiments, a single forwarding table is used to store forwarding entries associated with different tiers of traffic.

(57) In some embodiments, the generation of forwarding table entries for forwarding a flow of traffic to a remote PoP is dynamic. A forwarding policy is a forwarding relationship between a first PoP and a second PoP, where traffic is forwarded from the first PoP to the second PoP. When the amount of traffic associated with a forwarding relationship increases, an appropriate number of forwarding table entries is generated or updated to forward flows to the remote PoP. Alternatively, when the amount of traffic associated with the forwarding relationship is reduced, an appropriate number of forwarding table entries is updated for forwarding the flows to local compute servers of the PoP receiving the traffic instead of forwarding traffic to the remote PoP. Changes to the forwarding policy causes the update of forwarding table entries.

(58) In some embodiments, the generation of forwarding table entries for forwarding traffic between PoPs ensures that traffic forwarded from a first PoP to a second PoP is not forwarded from the second PoP to a third PoP. Respecting this condition allows to avoid traffic looping from one PoP and back to itself via another PoP and avoids unnecessary latency that would be introduced by repeated forwarding of the same traffic to multiple PoPs.

(59) In some embodiments, the forwarding table entries that are selected to be updated when there is a change in the amount of traffic to forward (reduced or increased) for a PoP, the forwarding table entry with the least-recent changes is selected. This avoids having the same forwarding table entries (and consequently the same flows of traffic) being modified over and over. The selection of the forwarding tables with the least recent changes results in less connection being disturbed and the system maintaining these connections for a longer period of time while enabling load balancing of traffic across multiple PoPs in the cloud computing infrastructure.

(60) Forwarding Traffic from One PoP to Another

(61) FIG. 1E illustrates a block diagram of an exemplary forwarding unit for forwarding traffic to another PoP according to one embodiment. The forwarding table **126A** includes one or more forwarding entries that are generated from forwarding policies **142A** for forwarding traffic from PoP **150A** to PoP **150B**. A forwarding table entry includes a packet identifier, e.g., ID_1 and ID 2, and an associated destination for the packet. The packet identifier identifies the flow of packets received from client devices such as client device **110A**.

(62) In one embodiment, the packet identifier is generated at the packet ID determiner **180** upon receipt of the packet. The packet identifier can be generated by hashing one or more of the fields that define a flow of traffic that is to be received at the PoP. Additionally or alternatively, the packet identifier can be generated by selecting a portion of the hashed value (e.g., last N bits) obtained from the fields of the packet. For example, the packet ID may be determined by hashing a five-tuple of fields of the packet including the protocol version (e.g., the version of the IP protocol (IPv4 or IPv6, etc.), the source address (e.g., source IP address) and source port number (e.g., source TCP port number), a destination address (e.g., destination IP address) and destination port number (e.g.,

destination TCP port number) of a packet.

(63) In some embodiments, different fields of a packet of a flow of traffic can be used to generate the identifier. The fields can be selected to enable a certain forwarding behavior for the compute server **120A** of the first PoP **150A**. In one exemplary implementation, the choice of the five-tuple described above for generation of the packet identifier enables the forwarding of packets of different transport layer connections (e.g., TCP connections) from a single client device **110A** towards multiple PoPs of the cloud computing infrastructure. For example, a first connection from client device **110A** is defined with a first tuple including first protocol version, first source address, first source port number and first destination address and first destination number and a second connection from the client device **110A** is defined with a second tuple including second protocol version, second source address, second source port number, second destination address and second destination number, where some of the first fields and the second fields are different (e.g., first and second port number are different) resulting into two different identifiers for packets of these two transport layer connections. The two IDs can be associated with two different destinations **114** resulting in packets of the two connections providing from the same client device **110A** to be forwarded to two different PoPs. In an alternative implementation, the identifier of the packet can be generated based on the source address (e.g., source IP address) and the destination address (e.g., destination IP address) without including the port numbers and/or the protocol version, resulting in packets of different transport layer connections from the same client device **110A** to be associated with the same packet ID. In this exemplary implementation, the packets received from different connections of the client device **110A** with the compute server **120A** are forwarded according to a same forwarding entry towards a same forwarding destination.

(64) The destination **114** can include a network address **163** to forward traffic to the PoP **150B**. The address can be an external IP address of a compute server **120B**. The address can be a virtual IP address of a group of compute servers (which include the compute server **120B**) of the PoP **150B**. The address can be an external IP address of an edge network device in the PoP **150B** that is to receive incoming traffic to the PoP. For example, the edge network device can include the forwarding unit **124B**. Additionally or alternatively, the edge network device can include a router of the PoP **150B**. In some embodiments, the edge network device can include a decapsulation unit **190B**. The destination **114** further includes an inter-PoP indicator **164**. The inter-PoP indicator **164** indicates whether the packet is to be forwarded to a PoP that is different from the current PoP processing the packet. The inter-PoP indicator **164** may include a first value (e.g., 1) to indicate that the forwarding is inter-PoP (towards another PoP) or a second value (e.g., 0) to indicate that the forwarding is intra-PoP (within the same PoP).

(65) In operation a packet **184A** is received at the PoP **150A**. The packet ID determiner **180A** determines an identifier ID_1 of the packet **184A**. In some embodiments, the packet **184B** is an IP/TCP packet that includes a TCP header, an IP header, and a payload. The TCP header includes a source and destination port numbers. The IP header includes a source and destination IP address, and a version of the IP protocol. The identifier ID_1 is used to access a forwarding table entry in the forwarding table **126A**. The identifier ID_1 is associated with a forwarding table entry that includes an IP address of the PoP **150B** and an inter-PoP indicator that indicates that the packet is to be forwarded to another PoP. Based on the inter-PoP indicator, the packet **184A** is encapsulated according to an encapsulation protocol at the encapsulation unit **185A** and the destination address retrieved from the forwarding table. In some embodiments, the encapsulation protocol can be Generic Routing Encapsulation (GRE) protocol. In some embodiments, the encapsulation protocol can be another protocol (e.g., Generic User Datagram Protocol (UDP) Encapsulation (GUE) or other). The packet **184A** is encapsulated to include an encapsulation header **186A**. The encapsulation header **186A** includes the IP address of the first PoP **150A** as a source address and a destination IP address of the second PoP **150B** (e.g., external IP address of the compute server **120B**, an external IP address of an edge network device in the PoP **150B**, or a virtual IP address,

etc.). In some embodiments, the encapsulation protocol can be GRE, GUE, or another encapsulation protocol. The encapsulated packet **188** is transmitted to the PoP **150B** through a tunnel (e.g., GRE tunnel) over the network **105**. The PoP **150** receives the encapsulated packet **188** and decapsulates the encapsulated packet **188** at the decapsulation unit **190B** to obtain the packet **184A**.

(66) The encapsulated traffic including the encapsulated packet **188** travels over the existing network **105**. In some embodiments, the encapsulated traffic travels over an external network formed at least in part by the PoPs of the cloud computing infrastructure. The destination address of an encapsulated packet is an external IP address of a network device in the PoP **150B**, a virtual IP address in the PoP **150B**, or an external IP address of an edge network device.

(67) Following the decapsulation of the encapsulated packet **188** at the POP **150B**, the packet **184A** is processed at the POP **150B** as if it were received from the client device **110A**. In some embodiments, the packet **184A** is processed at the forwarding unit **124B** prior to being transmitted to the compute server **120B** for processing according to the cloud computing service. In other embodiments, the packet **184A** is directly transmitted to the compute server **120B** to be processed according to the cloud computing service. In some embodiments, the packet **184A** is processed at the compute server **120B** with zero or more other packets of the same flow to obtain the request for the second resource (e.g., one or more packets including the packet **184A** can be processed at a network layer to obtain an application layer request (e.g., HTTP request) for the second request). The request for the second resource is processed by the service **123B**. In some embodiments, when the request for the second network resource is processed, the request is fulfilled by retrieving the second network resource (from a cache or by transmitting a request to an origin server and receiving a response from the origin server). The compute server **120B** transmits the second response to the client device **110A** in one or more packets (IP packets). In some embodiments, the packet **184A** may be forwarded to another PoP instead of being processed at the application layer in the POP **150B**.

(68) Forwarding of Traffic to Another PoP while Maintaining Existing Connections

(69) In some embodiments, the update made to the forwarding tables **122A** in response to the receipt of a new inter-PoP forwarding policy does not disrupt established connections between the client devices and the PoPs. That is, changes made to the inter-PoP forwarding policy of a PoP ensures that established connections are maintained such that traffic received through these connections continues to be received without interruption. The forwarding mechanisms described herein enable new connections to obey new inter-PoP forwarding policies while ensuring that existing connections follow previous forwarding policies without interruption. In some embodiments, the support of existing connections can be temporary such that a connection that is maintained can be discontinued after a period of time. In some embodiments, the period of time can be determined empirically by the central controller. In other embodiments, the period of time can be set by an administrator of the cloud computing infrastructure **114**.

(70) In one example, a first forwarding policy can be configured in the first PoP **150A** to cause a first flow of traffic received from client device **110A** through a first connection at the first PoP **150A** to be forwarded to the second PoP **150B** instead of being processed at the first PoP **150A**. An update of the first forwarding policy at a later time causes traffic received from the same client device through a new connection to be forwarded to another PoP (e.g., the local PoP or a different remote PoP).

(71) FIG. 1F illustrates a block diagram of an exemplary forwarding unit for forwarding traffic to another PoP according to one embodiment where existing connections are maintained. Similarly to the embodiments described above with reference to FIG. 1E, the forwarding unit **124A** includes a packet ID determiner **180A** and an encapsulation unit **185A**. The packet ID determiner **180A** and the encapsulation unit **185A** perform similar operations as the one described with reference to FIG. 1E. The forwarding unit **124A** further includes forwarding table(s) **136A**. In some embodiments,

the forwarding table(s) **136A** include some differences with respect to the forwarding table(s) **126A** that will be described in further detail herein.

(72) The forwarding table(s) **136A** includes one or more forwarding table entries. A forwarding table entry includes a packet identifier, e.g., ID **1**, ID **2**. The packet identifier is determined as described in further detail above with respect to the operations of the packet ID determiner **180A**. A forwarding table entry includes a primary destination **172** and a secondary destination **174**. A primary destination **172** includes a network address of a network device. In some embodiments, the network address is an external IP address of a compute server in the PoP. The secondary destination **174** includes a network address **173** of a network device, an inter-PoP indicator **175**, and a forwarding direction **177**. In some embodiments, the network address is an external IP address of the network device. For example, the network address is an external IP address for a first PoP or an external IP address for a second PoP. The external IP address of a PoP can be an IP address of a compute server in the PoP, an edge network device in the PoP, or a virtual IP address that is routed to any one or more of compute servers in the PoP. The inter-PoP indicator indicates whether the packet is to be forwarded to another PoP that is different from PoP **150A**. The inter-PoP indicator **164** may include a first value (e.g., 1) to indicate that the forwarding is inter-PoP (towards another PoP) or a second value (e.g., 0) to indicate that the forwarding is intra-PoP (within the same PoP). The forwarding direction specifies the disposition of packets that initiate new connections. The forwarding direction can indicate that such packets are to be forwarded away, so that new connections will be forwarded to another PoP referred to by the secondary destination. Alternatively it can indicate that such packets are to be forwarded back, meaning that such packets and their associated connections are handled by the compute server referred to by the primary destination.

(73) The operation of the forwarding unit **124A** will be described with respect to packet **184B**. In some embodiments, the packet **184B** is an IP/TCP packet that includes a TCP header, an IP header, and a payload. The TCP header includes a source and destination port numbers. The IP header includes a source and destination IP address, and a version of the IP protocol. The PoP **150A** receives the first packet **184B**. The first packet includes as a first destination address an anycast address advertised by the first PoP for reaching the cloud computing service. The PoP **150A** identifies based on a forwarding table entry and a first identifier ID_1 for a first flow of the first packet, a primary destination **172** and a secondary destination **174**. The primary destination includes an external network address of a local compute server in the PoP **150A**. For example, the primary destination **172** includes the IP address of the compute server **120A**. The secondary destination includes an external network address of a second compute server that is located in a second PoP **150B** that is different from the first PoP. The secondary destination further includes an inter-PoP indicator **175** (here value=1) that indicates that the packet can be subject to remote forwarding. The second destination further includes a forwarding direction “away” indicating that the packet can be forwarded away from the current PoP **150A**.

(74) The first packet is transmitted to a first compute server **120A** of the first PoP **150A**. Transmitting the first packet to the first compute server **120A** includes encapsulating the first packet according to an encapsulation protocol in an encapsulated packet and the second network address of the second PoP is included in a field of an encapsulation header of the encapsulated packet. In some embodiments, the encapsulation protocol can be GRE, GUE, or another encapsulation protocol. In other embodiments, the first packet is processed locally at the first PoP without the need for transmitting to another local compute server. The PoP **150A** determines whether a first connection is already established for the first flow locally. When the packet is transmitted to the first compute server **120A**, the determination of whether a connection already exists for the packet is performed at the compute server **120A**. For example, the PoP **150A** (which may be performed by the compute server **120A** or another network device in the PoP **150A**) may determine whether a TCP connection exists in a first compute server **120A** for the flow of the

packet that is received. In response to determining that a first connection is established at the first compute server for the first flow, the first packet is processed at the first PoP **150A** according to the cloud computing service. In some embodiments, processing the packet in the first PoP **150A** results in the PoP **150A** transmitting one or more packets including a response to the client device **110A**. The response is for the request received at the PoP **150A** with one or more packets that include the first packet. In this exemplary scenario, while the flow of the first packet is defined to be forwarded to a remote PoP in the forwarding table entry, the forwarding of packets is performed such that packets received through existing connections (e.g., TCP connections) are processed locally without being forwarded (i.e., without being disrupted).

(75) Alternatively, when the PoP **150A** determines that there is no connection established for the flow of the first packet in the first PoP, the PoP **150A** determines whether the packet is an initial packet of a new connection for the first flow. When the first packet is not the initial packet of the new connection for the first flow, the PoP **150A** transmits the packet to the second PoP to be processed in the second PoP **150B** according to the cloud computing service. When the packet is not a packet for a new connection (e.g., a SYN packet in TCP) and there are no open connections for that flow in the local PoP, this is an indication that the packet belongs to a flow associated with an open connection in the remote PoP and therefore the packet is forwarded to the PoP **150B** for processing. The transmission of the packet to the remote PoP includes encapsulating the packet **184B** to obtain an encapsulated packet **188B**. The packet is encapsulated to include an encapsulation header. The encapsulation header includes the IP address of the first PoP **150A** and a destination IP address of the second PoP **150B** (e.g., external IP address of the compute server **120B**, an external IP address of an edge network device in the PoP **150B**, or a virtual IP address, etc.). In some embodiments, the encapsulation protocol can be GRE, GUE, or another encapsulation protocol.

(76) Alternatively, in response to determining that the packet is the initial packet of the new connection for the flow (e.g., a SYN packet in TCP) the first packet **184B** is processed according to the forwarding direction (here away). When the forwarding direction is “away,” this indicates that the forwarding of this flow to the remote PoP is enabled and the packet **184B** can be forwarded towards the second PoP **150B** for processing. The transmission of the packet **184B** to the remote PoP includes encapsulating the packet to obtain an encapsulated packet **188B**. The packet is encapsulated to include an encapsulation header. The encapsulation header includes the IP address of the first PoP **150A** and a destination IP address of the second PoP **150B** (e.g., external IP address of the compute server **120B**, an external IP address of an edge network device in the PoP **150B**, or a virtual IP address, etc.). In some embodiments, the encapsulation protocol can be GRE, GUE, or another encapsulation protocol. Alternatively, when the forwarding direction is “back,” this is an indication that the forwarding of packets to a remote PoP is not enabled and therefore the packet is to be processed locally at the first compute server **120A** of the first PoP **150A**.

(77) The embodiments herein enable the forwarding of traffic to remote PoPs while maintaining existing connections. The forwarding table entries can be configured/updated with appropriate forwarding directions to indicate whether the forwarding is enabled or not. The use of the two destinations (primary and secondary) allows to forward traffic to the local compute server or a remote compute server of another PoP based on whether a connection is established for this traffic at a respective compute server. Regardless of whether the remote forwarding is enabled or not, established connections are maintained in both the remote PoP and the local PoP.

(78) Exemplary Operations:

(79) FIG. 2A illustrates a flow diagram of exemplary operations that can be performed for determining forwarding policies for a PoP in the cloud computing infrastructure in accordance with some embodiments.

(80) At operation **202**, server **112** receives from the PoPs of the cloud computing infrastructure **101** feedback related to the state of compute servers in the PoPs. The feedback of a compute server is

indicative of a state of the compute server at a given time. In some embodiments, the feedback can be a parameter indicative of a CPU load of the compute server. Alternatively or additionally, the feedback can be a parameter that measures the Quality of Service (QoS) provided by a service implemented on the compute server. The feedback can be a measure of latency and/or reliability of the service at the compute server.

(81) In some embodiments, the feedback is periodically received from the compute servers of the PoPs. For example, each one of the compute servers from the PoPs **150A-N** may be configured to push to the server **112** the feedback at a regular time interval. In an alternative embodiment, the server **112** may be configured to request feedback from the compute servers of the PoPs **150A-N**. In some embodiments, the feedback of a compute server can be either transmitted from the compute server to the server **112** or alternatively, the feedback may be gathered at the PoP management server from one or more compute servers, e.g., PoP management server **140A**, and transmitted from the PoP management server **140A** to the server **112**.

(82) The flow of operations moves to operation **204**. At operation **204**, the server **112** identifies, based on the received feedback, one or more first PoPs from which traffic needs to be forwarded to another PoP. For example, the server **112** may identify the PoP **150A**. In some embodiments, the identification of a PoP can be performed based on the feedback received for one or more of the compute servers of the PoP. For example, the identification of the PoP can be performed based on the CPU load of the compute servers of the PoP. In some embodiments, PoP **150A** can be identified as a PoP from which traffic needs to be forwarded to another PoP when at least one of the compute server's CPU load exceeds a threshold value. In other embodiments, PoP **150A** can be identified when the CPU load of each one of multiple compute servers (e.g., a number N of compute servers, a percentage of the compute servers, etc.) is greater than the threshold value. In other embodiments, PoP **150A** can be identified when a combined CPU load of multiple compute servers of the PoP **150A** exceeds a threshold value. Alternatively or additionally, the identification of the PoP can be performed based on other criteria such as QoS, or type of service. For example, traffic of a particular type of service (e.g., video service) may have stricter requirements for latency/reliability compared to other types of services (e.g., regular web page surfing, email, etc.). Thus, a PoP that is receiving the type of traffic with stricter requirements may not be identified as a PoP from which traffic is to be forwarded. In other embodiments, the PoP **150A** may be selected and the type of service may be used in the determination of forwarding policies and/or forwarding table entries for the particular PoP.

(83) The flow of operations moves to operation **206**. At operation **206**, the server **112** identifies, based on the received feedback, one or more second PoPs towards which traffic can be forwarded from another PoP. In some embodiments, a PoP can be identified, e.g., PoP **150B**, based on the feedback received from one or more of the compute servers of the PoP. The feedback received from the compute servers can be a parameter indicative of the CPU load. In some embodiments, the server **112** considers all of the compute servers **120B-Q** of the PoP **150B**. For example, the server **112** may determine that individual compute servers have a CPU load that is lower than a predetermined threshold value. In some embodiments, the server **112** may evaluate whether at least one of the compute servers in the PoP **150B** has a CPU load that is lower than the threshold value to identify the PoP **150B** as a PoP that can receive forwarded traffic. Alternatively, the server **112** may identify the PoP **150B** when at least a number of compute servers (a predetermined number of compute servers, or a percentage of the compute servers) have a CPU load that is lower than the threshold value. Alternatively or additionally, the server **112** may determine that a combination of the compute servers **120B-Q** (e.g., a subset or all of the compute servers **120B-Q**) have a combined CPU load that is lower than a predetermined threshold value. In these embodiments, the CPU load of some of the compute servers of the PoP **150B** may have a relatively high load but a total load across the compute servers in the PoP does not exceed a threshold value. Alternatively or additionally, the identification of the PoP can be performed based on other criteria such as QoS or

type of service. For example, traffic of a particular type of service (e.g., video service) may have stricter requirements for latency/reliability compared to other types of services (e.g., regular web page surfing, email, etc.). Thus, a PoP that is receiving the type of traffic with stricter requirements may not be identified as a PoP to which traffic is to be forwarded.

(84) The flow of operations moves to operation **208**. At operation **208**, the server **112** determines the inter-PoP forwarding policies based on the identified PoPs. The inter-PoP forwarding policies enable load balancing of traffic between the PoPs of the cloud computing infrastructure **101**. A forwarding policy for a PoP includes 1) an amount of traffic that needs to be forwarded from the PoP and 2) a destination network address of a PoP towards which traffic is to be forwarded. In some embodiments, the destination network address can be an external address (e.g., external IP address) of a compute server of the PoP, where the compute server hosts the cloud computing service. Alternatively or additionally, the destination network address can be an external address of a network device in the PoP that has other functionalities. For example, the network device can be a router or a forwarding unit in the PoP.

(85) The flow of operations moves to operation **210**. At operation **210**, the server **112** transmits the forwarding policies **142** to one or more PoPs of the cloud computing infrastructure **101**. For example, the server **112** transmits the forwarding policies to PoP **150A** for forwarding traffic to PoP **150B**. The configuration of the PoPs with the forwarding policies results in at least one PoP, e.g., PoP **150A**, forwarding a portion of the traffic received according to the anycast protocol to another PoP, e.g., **150B**. In some embodiments, the forwarding policies for a PoP may cause the PoP to forward traffic to multiple ones of the other PoPs in the cloud computing infrastructure **101**.

(86) The determination of the inter-PoP policies **142** is a dynamic mechanism that is performed periodically based on the feedback received for the compute servers of the PoPs. In some embodiments, changes to the forwarding policies can occur at a high frequency such as multiple times per hour. In some embodiments, the inter-PoP forwarding policy that is determined, can be a new inter-PoP forwarding policy (e.g., a policy to forward traffic from PoP **150A** to PoP **150B**), an update to an existing inter-PoP policy, or a cancellation of a previously existing inter-PoP forwarding policy. For example, when the central controller **114** determines that an existing inter-PoP forwarding policy is to be updated, it may transmit a request to modify the fraction of traffic that is forwarded for an established forwarding relationship.

(87) Once the inter-PoP forwarding policies **142** are transmitted from the server **112** to the PoP(s) **150A-N**, they are converted into forwarding table entries in the PoP(s). For example, a PoP management server **140A** receives and stores the inter-PoP policies **142A**. The inter-PoP forwarding policies **142A** are converted into forwarding table entries **143A**. In some embodiments, the forwarding table entries **143A** are transmitted to one or more of compute servers **120A-P** of the PoP **150A** to be stored in forwarding tables of the compute servers. In some embodiments, different forwarding entries can be generated for different compute servers of the same PoP. For example, if the anycast address associated with a first compute server **120A** is different from the anycast address associated with a second compute server **120B**, different traffic will reach the different compute servers based on these two different anycast addresses. In this embodiment, different forwarding entries are populated for the different compute servers. In other embodiments, the same forwarding table entries are forwarded to multiple compute servers in a PoP.

(88) In some embodiments, the generation of forwarding table entries for forwarding a flow of traffic to a remote PoP is dynamic. A forwarding policy is a forwarding relationship between a first PoP and a second PoP, where traffic is forwarded from the first PoP to the second PoP. When the amount of traffic associated with a forwarding relationship increases, an appropriate number of forwarding table entries is generated or updated to forward flows to the remote PoP. Alternatively, when the amount of traffic associated with the forwarding relationship is reduced, an appropriate number of forwarding table entries is updated for forwarding the flows to local compute servers of the PoP receiving the traffic instead of forwarding traffic to the remote PoP. Changes to the

forwarding policy causes the update of forwarding table entries.

(89) In some embodiments, the generation of forwarding table entries for forwarding traffic between PoPs ensures that traffic forwarded from a first PoP to a second PoP is not forwarded from the second PoP to a third PoP. Respecting this condition allows to avoid traffic looping from one PoP and back to itself via another PoP and avoids unnecessary latency that would be introduced by repeated forwarding of the same traffic to multiple PoPs.

(90) FIG. 2B illustrates a flow diagram of exemplary operations that can be performed for load balancing traffic across PoP(s) of a cloud computing infrastructure, in accordance with some embodiments. The operation of FIG. 2B are performed in a first PoP of multiple PoPs of the cloud computing infrastructure. For example, the operations can be performed in PoP **150A**. In some embodiments, the operations herein are performed in a forwarding unit **124A**.

(91) At operation **212**, the PoP **150A** receives a packet that includes as a destination address an anycast address advertised by the first PoP for reaching a cloud computing service. In some embodiments, the flow of operations moves to operation **214**. In other embodiments, operation **214** is optional and the flow of operations moves to operation **216**. At operation **214**, the PoP **150A** determines whether the packet is a security threat. For example, the PoP **150A** may determine whether the packet is part of a volumetric attack (e.g., DDoS) on the PoP **150A**. When it is determined that the packet is a security threat on the PoP **150A**, the packet is dropped (operation **215**). Alternatively, when it is determined that the packet is not a security threat on the PoP **150A**, the flow of operations moves to operation **216**.

(92) At operation **216**, the PoP **150A** determines, based on a forwarding table entry and an identifier for a flow of the packet, a destination address of a second PoP of the plurality of PoPs that is different from the first PoP. The PoP **150A** determines an identifier ID_1 of the packet **184A**. The identifier ID_1 is used to access a forwarding table entry in the forwarding table **126A** and determines based on the packet identifier a forwarding table entry that includes an IP address of the PoP **150B**. The forwarding table entry further includes an inter-PoP indicator that indicates that the packet is to be forwarded to another PoP that is different from PoP **150A**.

(93) The flow of operations moves to operation **218**. At operation **218**, based on the inter-PoP indicator, the PoP **150A** encapsulates the packet into an encapsulated packet that includes the external address of the second Pop as the destination address of the encapsulated packet. For example, the packet **184A** is encapsulated according to an encapsulation protocol at the encapsulation unit **185A** and the destination address retrieved from the forwarding table. The packet is encapsulated to include an encapsulation header. The encapsulation header includes the IP address of the first PoP **150A** and a destination IP address of the second PoP **150B** (e.g., external IP address of the compute server **120B**, an external IP address of an edge network device in the PoP **150B**, or a virtual IP address, etc.). In some embodiments, the encapsulation protocol can be GRE, GUE, or another encapsulation protocol.

(94) The flow of operations moves to operation **220**. At operation **220**, the PoP **150A** transmits the encapsulated packet to the second PoP instead of processing the packet locally according to the cloud computing service. The encapsulated packet **188** is transmitted to the PoP **150B** through an encapsulation tunnel (e.g., GRE tunnel) over the network **105**. The encapsulated traffic including the encapsulated packet **188** travels over the existing network **105**. The destination address of an encapsulated packet is an external IP address of a network device in the PoP **150B**.

(95) Following the decapsulation of the packet at the POP **150B**, the packet **184A** is processed at the POP **150B** as if it were received from the client device **110A**. In some embodiments, the packet **184A** is processed at the forwarding unit **124B** prior to being transmitted to the compute server **120B** for processing according to the cloud computing service. In other embodiments, the packet **184A** is directly transmitted to the compute server **120B** to be processed according to the cloud computing service. In some embodiments, the packet **184A** is processed at the compute server **120B** with zero or more other packets of the same flow to obtain the request for the second

resource (e.g., one or more packets including the packet **184A** can be processed at a network layer to obtain an application layer request (e.g., HTTP request) for the second request). The request for the second resource is processed by the service **123B**. In some embodiments, when the request for the second network resource is processed, the request is fulfilled by retrieving the second network resource (from a cache or by transmitting a request to an origin server and receiving a response from the origin server). The compute server **120B** transmits the second response to the client device **110A** in one or more packets (IP packets).

(96) FIG. 2C illustrates a flow diagram of exemplary operations that can be performed for forwarding traffic from one PoP to another PoP while maintaining existing connections. At operation **230**, the PoP **150A** receives a first packet. The first packet includes as a first destination address an anycast address advertised by the first PoP for reaching the cloud computing service. The first packet is received at the first PoP from multiple PoPs of the cloud computing infrastructure as a result of an anycast protocol. The flow of operations moves to operation **231**. At operation **231**, the PoP **150A** identifies based on a first forwarding table entry and a first identifier for a first flow of the first packet, a first network address of the first PoP and a second network address of a second PoP that is different from the first PoP. The identification of the first network address and the second network address, includes a lookup of the first forwarding table entry in the forwarding table based on the first identifier of the first flow of the packet to retrieve the first network address and the second network address. The first network address is an external IP address of a first network device in the PoP. The first network device can be an edge network device that includes a router and/or a forwarding unit as described above. Additionally or Alternatively, the first network device can be a compute server. The second network address is an external IP address of a second network device in the PoP. The second network device can be an edge network device that includes a router and/or a forwarding unit as described above. Additionally or Alternatively, the second network device can be a compute server.

(97) In some embodiments, the first packet is transmitted to a first compute server **120A** of the first PoP **150A**, at operation **232**. Transmitting the first packet to the first compute server **120A** includes encapsulating the first packet in an encapsulated packet and the second network address of the second PoP is included in a field of an encapsulated header of the encapsulated packet. In some embodiments, the encapsulation protocol can be GRE, GUE, or another encapsulation protocol. In other embodiments, the first packet is processed locally at the first PoP without the need for transmitting to another local compute server. The flow moves to operation **233**. At operation **233**, the PoP **150A** determines whether a first connection is already established for the first flow locally. When the packet is transmitted to the first compute server **120A**, the determination of whether a connection already exists for the packet is performed at the compute server **120A**. For example, the PoP **150A** may determine whether a TCP connection exists in a first compute server **120A** for the flow of the packet that is received. In response to determining that a first connection is established at the first compute server for the first flow, the flow of operations moves to operation **234**. At operation **234**, the first packet is processed at the first PoP according to the cloud computing service.

(98) Referring back to FIG. 2C, in response to determining that there is no connection established for the flow of the first packet in the first PoP, the flow of operations moves to operation **236**. At operation **236**, the PoP **150A** determines whether the packet is an initial packet of a new connection for the first flow.

(99) In response to determining that the first packet is not the initial packet of the new connection for the first flow, the flow of operations moves to operation **238**. At operations **238**, the PoP **150A** transmits the packet to the second PoP to be processed in the second PoP according to the cloud computing service. Alternatively, in response to determining that the packet is the initial packet of the new connection for the flow, the flow of operations moves to operation **240**. At operation **240**, the first packet is processed according to a forwarding direction. In some embodiments, the PoP

150A further determines a forwarding direction from the forwarding table entry for the first packet. When accessing the forwarding table entry based on the first packet identifier, the PoP **150A** further retrieves the forwarding direction for the packet. The forwarding direction can indicate that the packet is to be forwarded away (remote forwarding) from the PoP or back to the PoP (local forwarding). In other words, the forwarding direction indicates whether the packet is to be forwarded to another PoP or alternatively, whether the packet is to be processed locally at the PoP **150A**.

(100) In response to determining that the forwarding direction indicates a local forwarding, the flow of operations moves to operation **234**. At operation **234**, the first packet is processed at the first PoP, e.g., at the first compute server **120A** according to the cloud computing service.

(101) In response to determining that the forwarding direction indicates a remote forwarding, the flow of operations moves to operation **238**. At operation **238**, the packet is encapsulated to obtain an encapsulated packet. The packet is encapsulated to include an encapsulation header. The encapsulation header includes the IP address of the first PoP **150A** and a destination IP address of the second PoP **150B** (e.g., external IP address of the compute server **120B**, an external IP address of an edge network device in the PoP **150B**, or a virtual IP address, etc.). In some embodiments, the encapsulation protocol can be GRE, GUE, or another encapsulation protocol.

(102) The flow of operations moves to operation **239**. At operation **239**, the packet is transmitted to the second PoP to be processed at the second PoP according to the cloud computing service. The encapsulated packet is transmitted to the PoP **150B** through a tunnel (e.g., GRE tunnel) over the network **105**. The encapsulated traffic including the encapsulated packet travels over the existing network **105**.

(103) FIG. 2D illustrates a flow diagram of exemplary operations that can be performed when receiving a packet from another PoP, in accordance with some embodiments.

(104) At operation **250** an encapsulated packet is received at the second PoP **150B**. The packet is decapsulated according to the encapsulation protocol (e.g., GRE, GUE, or other) to retrieve the original packet, at operation **252**. The flow moves to operation **254**, at which the packet is processed according to the cloud computing service at the second PoP **150B**. Thus, instead of being processed at the PoP that receives the packet based on the anycast protocol, the packet is forwarded to another PoP, the second PoP, and processed there consequently reducing the load on the first PoP. In some embodiments, upon receipt of the original packet, this packet can be processed as if it was directly received from a client device and can be subject to similar operations as described with reference to FIGS. 2B-C. In these embodiments, the packet may be forwarded to one or more other PoP prior to being processed at the application level.

(105) Exemplary Electronic Devices

(106) Electronic Device and Machine-Readable Media

(107) One or more parts of the above embodiments may include software and/or a combination of software and hardware. An electronic device (also referred to as a computing device, computer, etc.) includes hardware and software, such as a set of one or more processors coupled to one or more machine-readable storage media (e.g., magnetic disks, optical disks, read only memory (ROM), Flash memory, phase change memory, solid state drives (SSDs)) to store code (which is composed of software instructions and which is sometimes referred to as computer program code or a computer program) for execution on the set of processors and/or to store data. For instance, an electronic device may include non-volatile memory (with slower read/write times, e.g., magnetic disks, optical disks, read only memory (ROM), Flash memory, phase change memory, SSDs) and volatile memory (e.g., dynamic random access memory (DRAM), static random access memory (SRAM)), where the non-volatile memory persists code/data even when the electronic device is turned off or when power is otherwise removed, and the electronic device copies that part of the code that is to be executed by the set of processors of that electronic device from the non-volatile memory into the volatile memory of that electronic device during operation because volatile

memory typically has faster read/write times. As another example, an electronic device may include a non-volatile memory (e.g., phase change memory) that persists code/data when the electronic device is turned off, and that has sufficiently fast read/write times such that, rather than copying the part of the code/data to be executed into volatile memory, the code/data may be provided directly to the set of processors (e.g., loaded into a cache of the set of processors); in other words, this non-volatile memory operates as both long term storage and main memory, and thus the electronic device may have no or only a small amount of volatile memory for main memory. In addition to storing code and/or data on machine-readable storage media, typical electronic devices can transmit code and/or data over one or more machine-readable transmission media (also called a carrier) (e.g., electrical, optical, radio, acoustical or other form of propagated signals—such as carrier waves, infrared signals). For instance, typical electronic devices also include a set of one or more physical network interface(s) to establish network connections (to transmit and/or receive code and/or data using propagating signals) with other electronic devices. Thus, an electronic device may store and transmit (internally and/or with other electronic devices over a network) code and/or data with one or more machine-readable media (also referred to as computer-readable media).

(108) Electronic devices are used for a variety of purposes. For example, an electronic device (sometimes referred to as a server electronic device) may execute code that cause it to operate as one or more servers used to provide a service to another electronic device(s) (sometimes referred to as a client electronic device, a client computing device, or a client device) that executes client software (sometimes referred to as client code or an end user client) to communicate with the service. The server and client electronic devices may be operated by users respectively in the roles of administrator (also known as an administrative user) and end user.

(109) FIG. 3 illustrates a block diagram for an exemplary data processing system **300** that may be used in some embodiments. One or more such data processing systems **300** may be utilized to implement the embodiments and operations described with respect to the server **112**, the routers, the forwarding units **124**, the compute servers **120**, the origin servers **130**, and/or the client devices **110**.

(110) The data processing system **300** is an electronic device that stores and transmits (internally and/or with other electronic devices over a network) code (which is composed of software instructions and which is sometimes referred to as computer program code or a computer program) and/or data using machine-readable media (also called computer-readable media), such as machine-readable storage media **310** (e.g., magnetic disks, optical disks, read only memory (ROM), flash memory devices, phase change memory) and machine-readable transmission media (also called a carrier) (e.g., electrical, optical, radio, acoustical or other form of propagated signals—such as carrier waves, infrared signals), which is coupled to the processing system **320**. The processing system **320** may be one or more processors and/or connected system components such as multiple connected chips. The depicted machine-readable storage media **310** may store program code **330** that, when executed by the processor(s) **320**, causes the data processing system **300** to perform the operations described herein with reference to the block diagrams and the flow diagrams of FIGS. **1A-2D**.

(111) The data processing system **300** also includes one or more network interfaces **340** (e.g., a wired and/or wireless interfaces) that allows the data processing system **300** to transmit data and receive data from other computing devices, typically across one or more networks (e.g., Local Area Networks (LANs), the Internet, etc.). The data processing system **300** may also include one or more input or output (“I/O”) components **350** such as a mouse, keypad, keyboard, a touch panel or a multi-touch input panel, camera, frame grabber, optical scanner, an audio input/output subsystem (which may include a microphone and/or a speaker), other known I/O devices or a combination of such I/O devices. Additional components, not shown, may also be part of the system **300**, and, in certain embodiments, fewer components than that shown may be part of the system **300**.

(112) While the flow diagrams in the figures show a particular order of operations performed by certain embodiments, it should be understood that such order is exemplary (e.g., alternative embodiments may perform the operations in a different order, combine certain operations, overlap certain operations, etc.).

(113) While the above description includes several exemplary embodiments, those skilled in the art will recognize that the invention is not limited to the embodiments described and can be practiced with modification and alteration within the spirit and scope of the appended claims. The description is thus illustrative instead of limiting.

Claims

1. A method of load balancing traffic between a plurality of Points of Presence (PoPs) of a cloud computing infrastructure that provides a cloud computing service to one or more users, the method comprising: receiving feedback data from the plurality of PoPs, the feedback data related to states of compute servers in the plurality of PoPs; identifying, based on the feedback data, a first POP of the plurality of PoPs from which traffic needs to be forwarded to another POP; identifying, based on the feedback data, a second POP of the plurality of PoPs towards which traffic can be forwarded from another POP; determining an inter-POP forwarding policy defining an amount of first traffic received by the first POP to forward to the second POP; transmitting the inter-POP forwarding policy to the plurality of PoPs; receiving, at the first POP, the inter-PoP forwarding policy; converting, at the first PoP, the inter-POP forwarding policy to forwarding table entries of a forwarding table used in the first POP, wherein each forwarding table entry includes a packet identifier that corresponds to a packet flow and a network address associated with the second POP, wherein the forwarding table entries are configured to direct a proportion of packet flows that align with the defined amount of traffic received by the first POP to be forwarded to the second POP; receiving, at the first POP, a first packet; determining a first packet identifier of the first packet that identifies a first packet flow; determining, based on a first forwarding table entry of the forwarding table for the first packet identifier, a first network address associated with the second PoP; and transmitting the first packet to the first network address associated with the second PoP.
2. The method of claim 1, wherein identifying, based on the feedback data, the first POP of the plurality of PoPs from which the traffic needs to be forwarded to another PoP comprises: determining that a CPU load of at least one compute server at the first POP exceeds a threshold value.
3. The method of claim 1, wherein identifying, based on the feedback data, the second POP of the of the plurality of PoPs towards which the traffic can be forwarded from another POP comprises: determining that a CPU load of one or more compute servers at the second PoP is lower than a threshold value.
4. The method of claim 1, wherein determining the inter-PoP forwarding policy defining the amount of the first traffic received by the first POP to forward to the second POP further comprises: specifying a first amount of the first traffic of a first traffic type to be forwarded from the first POP to the second POP, wherein the first traffic type is associated with a first cloud computing service provided by the cloud computing infrastructure; and specifying a second amount of the first traffic of a second traffic type to be forwarded from the first POP to the second POP, wherein the second traffic type is associated with a second cloud computing service provided by the cloud computing infrastructure.
5. The method of claim 4, further comprising: receiving updated feedback data from the plurality of PoPs; and modifying the specified first amount of the first traffic of the first traffic type to be forwarded from the first POP to the second POP and the specified second amount of the first traffic of the second traffic type to be forwarded from the first POP to the second PoP based on the updated feedback data.

6. The method of claim 1, further comprising: identifying, based on the feedback data, a third POP of the plurality of PoPs towards which traffic can be forwarded from another POP; specifying a first amount of the first traffic of a first traffic type to be forwarded from the first POP to the second POP, wherein the first traffic type is associated with a first cloud computing service provided by the cloud computing infrastructure; and specifying a second amount of the first traffic of a second traffic type to be forwarded from the first POP to the third POP, wherein the second traffic type is associated with a second cloud computing service provided by the cloud computing infrastructure.
7. The method of claim 1, wherein the feedback data for a compute server in the plurality of PoPs includes one or more of a first parameter indicating a CPU load of the compute server, a second parameter measuring a Quality of Service provided by a service implemented on the compute server, and a measure of latency of the service.
8. A system, comprising: a first non-transitory machine-readable storage medium that provides instructions that, if executed by a first processor of a central controller of a cloud computing infrastructure that provides a cloud computing service to one or more users, will cause said central controller to perform operations including: receiving feedback data from a plurality of Point of Presences (PoPs) of the cloud computing infrastructure, the feedback data related to states of compute servers in the plurality of PoPs; identifying, based on the feedback data, a first POP of the plurality of PoPs from which traffic needs to be forwarded to another POP; identifying, based on the feedback data, a second POP of the plurality of PoPs towards which traffic can be forwarded from another POP; determining an inter-POP forwarding policy defining an amount of first traffic received by the first POP to forward to the second POP; and transmitting the inter-PoP forwarding policy to the plurality of PoPs; a second non-transitory machine-readable storage medium that provides instructions that, if executed by one or more second processors of the first PoP will cause operations to be performed including: receiving, at the first POP, the inter-POP forwarding policy; converting, at the first POP, the inter-PoP forwarding policy to forwarding table entries of a forwarding table used in the first POP, wherein each forwarding table entry includes a packet identifier that corresponds to a packet flow and a network address associated with the second POP, wherein the forwarding table entries are configured to direct a proportion of packet flows that align with the defined amount of traffic received by the first POP to be forwarded to the second PoP; receiving, at the first POP, a first packet; determining a first packet identifier of the first packet that identifies a first packet flow; determining, based on a first forwarding table entry of the forwarding table for the first packet identifier, a first network address associated with the second POP; and transmitting the first packet to the first network address associated with the second POP.
9. The system of claim 8, wherein identifying, based on the feedback data, the first POP of the plurality of PoPs from which the traffic needs to be forwarded to another POP further causes the central controller to perform operations including: determining that a CPU load of at least one compute server at the first POP exceeds a threshold value.
10. The system of claim 8, wherein identifying, based on the feedback data, the second POP of the plurality of PoPs towards which the traffic can be forwarded from another POP further causes the central controller to perform operations including: determining that a CPU load of one or more compute servers at the second POP is lower than a threshold value.
11. The system of claim 8, wherein determining the inter-POP forwarding policy defining the amount of the first traffic received by the first PoP to forward to the second POP further causes the central controller to perform operations including: specifying a first amount of the first traffic of a first traffic type to be forwarded from the first PoP to the second POP, wherein the first traffic type is associated with a first cloud computing service provided by the cloud computing infrastructure; and specifying a second amount of the first traffic of a second traffic type to be forwarded from the first POP to the second POP, wherein the second traffic type is associated with a second cloud computing service provided by the cloud computing infrastructure.
12. The system of claim 11, wherein the instructions further cause the central controller to perform

operations including: receiving updated feedback data from the plurality of PoPs; and modifying the specified first amount of the first traffic of the first traffic type to be forwarded from the first POP to the second PoP and the specified second amount of the first traffic of the second traffic type to be forwarded from the first PoP to the second PoP based on the updated feedback data.

13. The system of claim 8, wherein the instructions further cause the central controller to perform operations including: identifying, based on the feedback data, a third POP of the plurality of PoPs towards which traffic can be forwarded from another POP; specifying a first amount of the first traffic of a first traffic type to be forwarded from the first PoP to the second POP, wherein the first traffic type is associated with a first cloud computing service provided by the cloud computing infrastructure; and specifying a second amount of the first traffic of a second traffic type to be forwarded from the first POP to the third POP, wherein the second traffic type is associated with a second cloud computing service provided by the cloud computing infrastructure.

14. The system of claim 8, wherein the feedback data for a compute server in the plurality of PoPs includes one or more of a first parameter indicating a CPU load of the compute server, a second parameter measuring a Quality of Service provided by a service implemented on the compute server, and a measure of latency of the service.

15. A system, comprising: a central controller of a cloud computing infrastructure that includes a first processor; and a non-transitory machine-readable storage medium that provides instructions that, if executed by the first processor, will cause the central controller to perform operations including: receiving feedback data from a plurality of Point of Presences (PoPs) of the cloud computing infrastructure, the feedback data related to states of compute servers in the plurality of PoPs; identifying, based on the feedback data, a first POP of the plurality of PoPs from which traffic needs to be forwarded to another PoP; identifying, based on the feedback data, a second POP of the plurality of PoPs towards which traffic can be forwarded from another POP; determining an inter-POP forwarding policy defining an amount of first traffic received by the first POP to forward to the second POP; and transmitting the inter-POP forwarding policy to the plurality of PoPs; the first POP that includes a plurality of compute servers that each include a second non-transitory machine readable storage medium that provides instructions that, if executed by one or more second processors of the first POP will cause operations to be performed including: receiving, at the first POP, the inter-PoP forwarding policy; converting, at the first POP, the inter-POP forwarding policy to forwarding table entries of a forwarding table used in the first POP, wherein each forwarding table entry includes a packet identifier that corresponds to a packet flow and a network address associated with the second PoP, wherein the forwarding table entries are configured to direct a proportion of packet flows that align with the defined amount of traffic received by the first POP to be forwarded to the second POP; receiving, at the first POP, a first packet; determining a first packet identifier of the first packet that identifies a first packet flow; determining, based on a first forwarding table entry of the forwarding table for the first packet identifier, a first network address associated with the second POP; and transmitting the first packet to the first network address associated with the second PoP.

16. The system of claim 15, wherein identifying, based on the feedback data, the first POP of the plurality of PoPs from which the traffic needs to be forwarded to another POP causes the central controller to perform operations comprising: determining that a CPU load of at least one compute server at the first POP exceeds a threshold value.

17. The system of claim 15, wherein identifying, based on the feedback data, the second POP of the plurality of PoPs towards which the traffic can be forwarded from another POP causes the central controller to perform operations comprising: determining that a CPU load of one or more compute servers at the second POP is lower than a threshold value.

18. The system of claim 15, wherein determining the inter-PoP forwarding policy defining the amount of the first traffic received by the first POP to forward to the second POP causes the central controller to perform operations comprising: specifying a first amount of the first traffic of a first

traffic type to be forwarded from the first PoP to the second POP, wherein the first traffic type is associated with a first cloud computing service provided by the cloud computing infrastructure; and specifying a second amount of the first traffic of a second traffic type to be forwarded from the first POP to the second POP, wherein the second traffic type is associated with a second cloud computing service provided by the cloud computing infrastructure.

19. The system of claim 18, wherein the operations performed by the central controller further comprise: receiving updated feedback data from the plurality of PoPs; and modifying the specified first amount of the first traffic of the first traffic type to be forwarded from the first POP to the second POP and the specified second amount of the first traffic of the second traffic type to be forwarded from the first POP to the second PoP based on the updated feedback data.

20. The system of claim 15, wherein the operations performed by the central controller further comprise: identifying, based on the feedback data, a third POP of the plurality of PoPs towards which traffic can be forwarded from another POP; specifying a first amount of the first traffic of a first traffic type to be forwarded from the first PoP to the second POP, wherein the first traffic type is associated with a first cloud computing service provided by the cloud computing infrastructure; and specifying a second amount of the first traffic of a second traffic type to be forwarded from the first POP to the third POP, wherein the second traffic type is associated with a second cloud computing service provided by the cloud computing infrastructure.

21. The system of claim 15, wherein the feedback data for a compute server in the plurality of PoPs includes one or more of a first parameter indicating a CPU load of the compute server, a second parameter measuring a Quality of Service provided by a service implemented on the compute server, and a measure of latency of the service.
