

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12395430
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	De Grace; Gerald Roy et al.

Disaggregation of TIER1 devices in an SDN using SmartSwitches

Abstract

Techniques are disclosed for processing data packets and implementing policies in a software defined network (SDN) of a virtual computing environment. At least one network device is configured to disaggregate enforcement of policies of the SDN from hosts of the virtual computing environment. Tier-0 devices are communicatively coupled to network interfaces of the network device. The network device comprises a plurality of data processing units that are configured to implement functionality of the network device.

Inventors: De Grace; Gerald Roy (Atlanta, GA), Grantham; James Allen (Woodinville, WA), Moore; Kristina E. (Redmond, WA), Yuan; Lihua (Redmond, WA), Ge; Nan (Redmond, WA)

Applicant: MICROSOFT TECHNOLOGY LICENSING, LLC (Redmond, WA)

Family ID: 1000008768066

Assignee: MICROSOFT TECHNOLOGY LICENSING, LLC (Redmond, WA)

Appl. No.: 17/855675

Filed: June 30, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20230370377 A1	Nov. 16, 2023

Related U.S. Application Data

us-provisional-application US 63342044 20220513

Publication Classification

Int. Cl.: H04L45/76 (20220101); H04L49/1515 (20220101)

U.S. Cl.:

CPC H04L45/76 (20220501); H04L49/1515 (20130101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
9014191	12/2014	Mandal	370/392	H04L 45/04
2015/0023368	12/2014	Connolly	370/420	H04L 12/6402
2016/0373359	12/2015	Bruun	N/A	H04L 45/64

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
2911347	12/2014	EP	N/A

OTHER PUBLICATIONS

“Agilio OVS Software Architecture—White Paper”, Retrieved From:

<http://web.archive.org/web/20160428150109/>

https://www.netronome.com/media/redactor_files/WP_Agilio_SW.pdf, Apr. 28, 2016, 6 Pages.
cited by applicant

“International Search Report and Written Opinion Issued in PCT Application No.

PCT/US2023/019010”, Mailed Date: Jun. 29, 2023, 11 Pages. (MS# 411732-WO-PCT) cited by applicant

Pfaff, et al., “The Design and Implementation of Open vSwitch”, In Proceedings of the 12th USENIX Symposium on Networked Systems Design and Implementation, May 4, 2015, pp. 117-130. cited by applicant

Primary Examiner: Orgad; Edan

Assistant Examiner: Rivas; Raul

Attorney, Agent or Firm: Newport IP, LLC

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION (1) This non-provisional utility application claims priority to U.S. Patent Application Ser. No. 63/342,044 entitled “DISAGGREGATION OF TIER1 DEVICES IN AN SDN USING SMARTSWITCHES” and filed on May 13, 2022, which is hereby incorporated in its entirety by reference.

BACKGROUND

(1) A data center may house computer systems and various networking, storage, and other related

components. Data centers may, for example, be used by service providers to provide computing services to businesses and individuals as a remote computing service or provide “software as a service” (e.g., cloud computing). Software defined networking (SDN) enables centralized configuration and management of physical and virtual network devices as well as dynamic and scalable implementation of network policies. The efficient processing of data traffic is important for maintaining scalability and efficient operation in such networks.

(2) It is with respect to these considerations and others that the disclosure made herein is presented.

SUMMARY

(3) One architecture for implementing cloud-based computing includes connection of a plurality of servers to what is typically referred to as top-of-rack switch or ToR. In practice, the ToR may be placed in the middle of rack in order to shorten the cable lengths. However, the functionality of the ToR remains the same, which is to provide an in-rack network switching capability.

(4) From this point, the ToR may be connected to a plurality of Tier1 switches in a Clos configuration. If there are N Tier1 switches, then each ToR will connect at least once to all N Tier1 switches. The Tier1 may then connected in a Clos configuration to some number of larger Tier2 switches, typically covering a datacenter's worth of servers, which are then connected to even larger Tier3 switches to form a metro network. Many cloud networks use some form of Clos configuration to build out their datacenter and metro connectivity.

(5) A SmartToR, which may generally be referred to as a SmartSwitch, may be a switch that includes functionality of one or more SmartNICs. As used herein, a SmartNIC may be a hardware-based acceleration device that may implement various ways of leveraging hardware acceleration and offloading techniques to perform a function, such as, for example, implementing tasks in hard ASIC logic, implementing tasks in soft (configurable) FPGA logic, implementing some tasks as software on FPGA software processor overlays, implementing some tasks as software on hard ASIC processors, or a combination thereof. In some embodiments, the hardware-based acceleration device may be a network communications device, such as a network interface card (NIC). Such a NIC may be referred to herein as a SmartNIC.

(6) In one embodiment of a SmartSwitch implementation, one or more SmartSwitches may serve a plurality of servers to provide high performance or enhanced SDN networking. Additionally, virtual machines (VMs) may be dynamically associated to a SmartSwitch. A SmartSwitch may be a network switch with programmable DPUs, IPU, EPUs, etc. with SmartNIC hardware integrated into the switch design, with a selected number of options for management. The installation of a SmartSwitch at one or more Tier1 positions in the data center architecture enable more efficient utilization of computing and networking resources. For example, the described techniques can allow for virtual computing environments to support a variety of configurations including custom hardware and hybrid architectures while maintaining efficient use of computing resources such as processor cycles, memory, network bandwidth, and power.

(7) This Summary is provided to introduce a selection of concepts in a simplified form that are further described below in the Detailed Description. This Summary is not intended to limit the scope of the claimed subject matter. Furthermore, the claimed subject matter is not limited to implementations that solve any or all disadvantages noted in any part of this disclosure.

Description

DRAWINGS

(1) The Detailed Description is described with reference to the accompanying figures. In the description detailed herein, references are made to the accompanying drawings that form a part hereof, and that show, by way of illustration, specific embodiments or examples. The drawings herein are not drawn to scale. Like numerals represent like elements throughout the several figures.

- (2) FIG. 1A is a diagram illustrating an example architecture in accordance with the present disclosure;
- (3) FIG. 1B is a diagram illustrating an example architecture in accordance with the present disclosure;
- (4) FIG. 1C is a diagram illustrating an example architecture in accordance with the present disclosure;
- (5) FIG. 1D is a diagram illustrating an example architecture in accordance with the present disclosure;
- (6) FIG. 1E is a diagram illustrating an example architecture in accordance with the present disclosure;
- (7) FIG. 2A is a diagram illustrating an example architecture in accordance with the present disclosure;
- (8) FIG. 2B is a diagram illustrating an example architecture in accordance with the present disclosure;
- (9) FIG. 2C is a diagram illustrating an example architecture in accordance with the present disclosure;
- (10) FIG. 2D is a diagram illustrating an example architecture in accordance with the present disclosure;
- (11) FIG. 2E is a diagram illustrating an example architecture in accordance with the present disclosure;
- (12) FIG. 2F is a diagram illustrating an example architecture in accordance with the present disclosure;
- (13) FIG. 3 is a diagram illustrating an example architecture in accordance with the present disclosure;
- (14) FIG. 4 is a diagram illustrating an example architecture in accordance with the present disclosure;
- (15) FIG. 5 is a diagram illustrating an example architecture in accordance with the present disclosure;
- (16) FIG. 6A is a flowchart depicting an example procedure in accordance with the present disclosure;
- (17) FIG. 6B is a flowchart depicting an example procedure in accordance with the present disclosure;
- (18) FIG. 7 is a diagram illustrating a data center in accordance with the present disclosure;
- (19) FIG. 8 is a diagram illustrating a data center in accordance with the present disclosure;
- (20) FIG. 9 is an example computing system in accordance with the present disclosure;

DETAILED DESCRIPTION

(21) The disclosed embodiments enable datacenters to provide services in a manner that can enhance system flexibility and efficiency while reducing cost and complexity, allowing for more efficient use of computing, storage, and network resources. Efficient implementation of the end-to-end services by a cloud service provider can enable an experience that is seamless and more consistent across various footprints. The effective and efficient distribution of the described disaggregation and pooling techniques can be determined based on the implications for various performance and security implications such as latency and data security.

(22) In an embodiment, a ToR switch may be enabled with SmartNICs. Such a ToR switch may be referred to as a SmartSwitch or a SmartToR. In such a SmartSwitch configuration, most if not all software defined networking (SDN) may be performed at the SmartSwitch, allowing for the servers to deploy standard NICs (NICs which do not have SmartNIC capability—which may also be referred to herein as “skinny NIC”) which perform a lesser amount, if any, SDN functionality. In the disclosed embodiments, the SmartNIC functionality may be implemented in the Tier1 switches. This may be especially useful when only a fraction of the VMs of a cluster of servers require

disaggregated processing of a SmartSwitch or SmartToR.

(23) In some embodiments, no changes need to be implemented to the Clos networking design. If the SmartSwitch acts as the Tier1, then there will be no need to add the connectivity to an external SmartSwitch or other Smart Appliance. This can allow for an efficient design by removing interconnecting links or network hops that would otherwise be required to wire up a separate appliance. The number of data processing units (DPUs) in a SmartSwitch may be selected to match the expected load of the cluster of servers below the Tier1 device and can vary from one Tier1 design to another.

(24) The disclosed embodiments may be advantageously utilized when the Tier1 device is moved to chassis-based designs. In one embodiment, DPUs in the chassis-based design may be placed into DPU cards. In an embodiment, a DPU card may comprise several DPUs.

(25) Tier1 devices are typically deployed in sets of greater than 4 for redundancy, high availability, and resiliency. This may allow for improved high availability (HA) as compared to a Tier0 approach, where greater redundancy can be achieved in the event that one of the SmartSwitches fails.

(26) Additionally, the bandwidth to run a cluster and the DPUs stays within the Tier1 domain and hence does not require extra Tier2 links. This can be advantageous because every Tier2 link that is consumed reduces the number of server clusters that can be deployed.

(27) Referring to the appended drawings, in which like numerals represent like elements throughout the several FIGURES, aspects of various technologies for network disaggregation techniques and supporting technologies will be described. In the following detailed description, references are made to the accompanying drawings that form a part hereof, and which are shown by way of illustration specific configurations or examples.

(28) In the illustrated example scenarios, SDN capabilities may be enhanced by disaggregating policy enforcement from hosts and moving it onto a SmartSwitch strategically placed in the network. Software defined networking (SDN) is conventionally implemented on a general-purpose compute node. The SDN control plane may program the host to provide core network functions such as security, virtual network, and load balancer policies.

(29) The disclosed technologies may be part of a set of systems and technologies that aim to improve network performance to cloud applications through implementation of APIs and object models describing network services for the cloud. Such technologies may enable optimization of network performance and hardware/software technology to improve stateful connection performance. Example applications include NIC on a host, a SmartSwitch as disclosed herein, network disaggregation, and high-performance network appliances.

(30) In some implementations that use a rack level switch such as a top-of-rack (ToR) switch and higher tier network switches, the disclosed SmartSwitch may incorporate the capability to perform data traffic transforms and may be placed at various Tier1 locations in the network. The SmartSwitch can be used to provide transformations and connectivity. The SmartSwitch can accept policies that perform packet transformations. The traffic sent by workloads can be directed through the SmartSwitch, which can apply policies and perform transformations on the traffic and send the traffic to their destination.

(31) Since SmartSwitches can become a single point of failure for software defined networks, mitigation of faults for SmartSwitches may take into account the preservation of transient states (for example, TCP flow state) as well as the locality of the state within the individual SmartSwitches.

(32) The described embodiments may support, for example, connected devices such as FPGAs on SmartSwitches in multiple different network and physical topologies.

(33) The various aspects of the disclosure are described herein with regard to certain examples and embodiments, which are intended to illustrate but not to limit the disclosure. It should be appreciated that the subject matter presented herein may be implemented as a computer process, a

computer-controlled apparatus, a computing system, an article of manufacture, such as a computer-readable storage medium, or a component including hardware logic for implementing functions, such as a field-programmable gate array (FPGA) device, a massively parallel processor array (MPPA) device, a graphics processing unit (GPU), an application-specific integrated circuit (ASIC), a multiprocessor System-on-Chip (MPSoC), etc.

(34) A component may also encompass other ways of leveraging a device to perform a function, such as, for example, a) a case in which at least some tasks are implemented in hard ASIC logic or the like; b) a case in which at least some tasks are implemented in soft (configurable) FPGA logic or the like; c) a case in which at least some tasks run as software on FPGA software processor overlays or the like; d) a case in which at least some tasks run as software on hard ASIC processors or the like, etc., or any combination thereof. A component may represent a homogeneous collection of hardware acceleration devices, such as, for example, FPGA devices. On the other hand, a component may represent a heterogeneous collection of different types of hardware acceleration devices including different types of FPGA devices having different respective processing capabilities and architectures, a mixture of FPGA devices and other types hardware acceleration devices, etc.

(35) The functionality of the DPU engines within a SmartSwitch may provide a bump-in-wire where SDN policies and transformations for the cloud may be implemented. The bump, which may refer to a hop in a network path, does not add significant latency for packets that would have traveled in the same path due to routing. For example, if a packet stays within a compute row, the packet normally travels through the Tier1 switching complex to get from one rack to the other. Also, if the packet is to leave the cluster, the packet would likewise have to traverse the Tier1 switch. The actual latency of a DPU may be around 2 usec or less in practice, whether it is located on the server or the switch.

(36) SmartSwitches may be designed to handle VMs that require a high amount of connection performance. A typical Tier1 device with 1.6 Tbps of DPU each can handle all of these high performant VMs and have spare capacity. The actual capacity of the DPUs may change depending on the desired performance. If it is desired to change an existing deployment that was not implemented with high performant DPU technology to incorporate a SmartSwitch, there would thus be little additional delay other than the aforementioned latency to travel to the SmartSwitch and back.

(37) New cluster types may be identified to further deploy SmartSwitches to create new services. Some examples of cluster types may include graphics processing unit (GPU), artificial intelligence (AI), high performance compute (HPC), or bare metal compute clusters. Such clusters may benefit from having their own dedicated SmartSwitch complex in the Tier1/Tier0 architecture. Each of these cluster types can benefit from high performant virtualization enabled by the technologies disclosed herein. By using SmartSwitches at the Tier1/Tier0 part of the network, any VM in the fleet that access these high performance compute applications can benefit. For example, a GPU cluster can now be virtualized such that the transactions per second entering/leaving the GPU cluster per VM can be programmed from hundreds of thousands of transactions per second to even millions of transactions per second. This creates new opportunities for services that may be referred to as virtualized GPU services. Similar opportunities may be enabled for virtualized AI, virtualized high speed compute, virtualized bare metal, etc., thus enabling different business units to optimize performance for their particular technologies.

(38) In all of these cases, the cluster inherently designed for high transaction rates can highly benefit from SmartSwitches that enable any VM in the fleet to simply bypass all SDN network processing on a native server and instead be tunneled to the dedicated SmartSwitches in the cluster. The SmartSwitches in the cluster may in turn process connections and associated packets that enter/leave the cluster. This provides flexibility for existing VMs to enter into these new services and extend their life for years beyond what may have been planned.

(39) When DPUs are in line with normal networking paths, there is no added latency as compared to providing the same processing on the server. This allows the bump-wire-processing to be placed in the tiers of a cluster or in fact, in the Tier0/Tier1 of another cluster that dedicates the high speed SDN processing to tasks that run on its own cluster.

(40) With reference to FIG. 1A, illustrated is an example of typical processing flows to a network and from a network. Traffic from a network may be processed with flow tables **105** and forwarded to their destination. The forwarded traffic is received at physical NICs **106** at their destination, where the traffic may be processed by virtual switches **107**. The processed traffic may be sent to virtual interfaces at a host VM **108**. The reverse process is shown in operations **101**, **102**, **103**, and **104**. FIG. 1B illustrates the processing shown in FIG. 1A in a bidirectional view.

(41) With reference to FIG. 1C, offload of functionality **120** to accelerator hardware may include processing **122** of GFTs and processing performed by a VFP and input/output processing by the elastic network interface (ENI) **121**.

(42) FIG. 1D illustrates an example of a SmartSwitch **130** with a switch ASIC **131** configured to perform offload processing **132**. A number of offloads may be configured. The example SmartSwitch **130** includes one example implementation with a number of 400 Gb uplinks, and a number of connections to racks for connecting to servers.

(43) FIG. 1E illustrates an example implementation of the use of logical connections, such as a VXLAN tunnel, that allows for data traffic to bypass processing on the host. The ENI at host VM **143** may have a VXLAN Tunnel A **147** starting from the host VM **143** that bypasses the Vswitch/NFP **142**, physical NIC **141**, and GFT **140**. The VXLAN Tunnel B **148** may bypass ENI **144** and VFP/GFT **145** at the remote host, and end at the ENI **146** at the remote host.

(44) FIG. 2A illustrates an example architecture where a plurality of servers **204** are connected to ToRs **203** in each rack. The ToRs **203** may be connected to a plurality of Tier1 switches **202** in a Clos configuration. If there are 8 Tier1 **202** switches, then each ToR **203** will connect at least once to all 8 Tier1 **202** switches. The Tier1 switches **202** may then connected in a Clos configuration to some number of Tier2 switches **201**.

(45) FIG. 2B illustrates an example where the Tier1 **202** switches are replaced with SmartSwitches **205**. In some implementations, Tier1 may also be referred to as Middle-of-Rack (MoR). The servers thus each have access to SmartSwitches **205** while bypassing processing at the local host. In typical scenarios, most VM traffic do not require offloaded tasks for processing by SmartSwitches **205** and are processed locally. Thus only VM traffic that requires high connections per second (CPS) need be sent directly to the SmartSwitches **205**.

(46) FIG. 2C illustrates an example where the SmartSwitches **205** enable disaggregated GPU virtualization in a cluster, illustrating that new cluster types may be identified to further deploy SmartSwitches to create new services. The GPU cluster may benefit from having its own dedicated SmartSwitch cluster **206** in the Tier1/Tier0 architecture. By locating SmartSwitches **205** at the Tier1/Tier0 part of the network in the cluster, any VM in the compute cluster **207** can access SmartSwitch cluster **206**.

(47) FIG. 2D illustrates an example where the SmartSwitches **205** enable AI disaggregation in a manner similar to GPU disaggregation in FIG. 2C. FIGS. 2E and 2F illustrates an example where the SmartSwitches **205** enable bare metal disaggregation. FIG. 2F further illustrates that traffic from hosts in compute cluster **207** may be tunneled to interfaces at ENI **144**, allowing for offloaded packet processing to be performed by SmartSwitches **205**.

(48) The example of FIG. 2F further illustrates the advantages of placing the SmartSwitches **205** at the Tier1 level. As shown, each port on a Tier 2 switch represents connectivity to an entire compute cluster, whereas such a blast radius per port does not exist at the Tier1 level. However, placing the SmartSwitches **205** at the Tier0 level, would result in overcapacity. Thus, Tier1 placement allows for an efficient balance between capacity and fault tolerance.

(49) With reference to FIG. 3, illustrated is an example of a rack **300** with two SmartSwitches **310**

having one or more smart NICs **320** and a plurality of compute rows **330** having servers. Any virtual machine **340** running on any server in the data center smart rack can utilize the SmartSwitches **310**. For example, virtual machines with a high connections per second (CPS) or flow scale needs can send flows through the SmartSwitches. The SmartSwitches may be configured to perform SDN data path functions at a significantly faster rate as compared to conventional methods.

(50) With reference to FIG. **4**, illustrated is an example of SDN disaggregation where non-compute functionality is removed off the compute host. In one implementation, smartNICs **410** may be included in SmartSwitches **400**. In an embodiment, the SmartSwitches **400** may be cost optimized. The SmartSwitches **400** may be configured to perform all SDN data path functions. In this and other figures herein, the dashed line indicates skinny NICs **440** and the solid line indicates smartNICs **410**. The skinny NICs **440** may be implemented on the servers **430** for low function, cost, and low power.

(51) FIG. **4** illustrates an example of a fault tolerant scheme that is resilient to a single failure. Each SmartSwitch **400** is cross-connected to each Tier0 ToR **420**. In an embodiment, at least two SmartSwitches **400** may use connection state replication. In an embodiment, at least two Tier0 ToRs **420** may use connection state replication. In one example, 2-4 SmartNICs **440** per SmartSwitch may be implemented depending on the load. In an embodiment, each ToR **420** may be fully connected to every SmartSwitch **400**. For example, each SmartSwitch **400** may provide two redundant 40 G bump-in-wire SDNs. The servers **430** may have dual ported skinny NICs **440**. In an embodiment, SDN agents **450** may execute on the SmartSwitches **400**. In an embodiment, one to eight SmartSwitches **400** may be implemented at the Tier1 level.

(52) With reference to FIG. **5**, the left side of the figure shows a skinny or standard NIC **520** that provides RDMA offload functionality **525** and connectivity to the appliance. Networking traffic may be tunneled to the appliance **530**. The right side of the figure shows stateful network policy-based forwarding and security **550** performed on SmartNICs **540**.

(53) Turning now to FIG. **6A**, illustrated is an example operational procedure for processing data packets and implementing policies in a software defined network (SDN) of a virtual computing environment, by at least one SmartSwitch network device configured to disaggregate enforcement of policies of the SDN from hosts of the virtual computing environment. In an embodiment, the hosts may be implemented on servers communicatively coupled to network interfaces of the SmartSwitch network device. In an embodiment, the servers may host a plurality of virtual machines. In an embodiment, the SmartSwitch network device comprises a plurality of smart network interface cards (sNICs) configured to implement functionality of the SmartSwitch network device. In an embodiment, the SmartSwitch network device is situated at a tier-1 position in the virtual computing environment. Such an operational procedure can be provided by one or more components illustrated in FIGS. **1** through **5**. The operational procedure may be implemented in a system comprising one or more computing devices. It should be understood by those of ordinary skill in the art that the operations of the methods disclosed herein are not necessarily presented in any particular order and that performance of some or all of the operations in an alternative order(s) is possible and is contemplated. The operations have been presented in the demonstrated order for ease of description and illustration. Operations may be added, omitted, performed together, and/or performed simultaneously, without departing from the scope of the appended claims.

(54) It should also be understood that the illustrated methods can end at any time and need not be performed in their entirety. Some or all operations of the methods, and/or substantially equivalent operations, can be performed by execution of computer-readable instructions included on a computer-storage media, as defined herein. The term “computer-readable instructions,” and variants thereof, as used in the description and claims, is used expansively herein to include routines, applications, application modules, program modules, programs, components, data structures, algorithms, and the like. Computer-readable instructions can be implemented on various

system configurations, including single-processor or multiprocessor systems, minicomputers, mainframe computers, personal computers, hand-held computing devices, microprocessor-based, programmable consumer electronics, combinations thereof, and the like.

(55) It should be appreciated that the logical operations described herein are implemented (1) as a sequence of computer implemented acts or program modules running on a computing system such as those described herein) and/or (2) as interconnected machine logic circuits or circuit modules within the computing system. The implementation is a matter of choice dependent on the performance and other requirements of the computing system. Accordingly, the logical operations may be implemented in software, in firmware, in special purpose digital logic, and any combination thereof. Thus, although the routine **600** is described as running on a system, it can be appreciated that the routine **600** and other operations described herein can be executed on an individual computing device or several devices.

(56) Referring to FIG. **6**, operation **601** illustrates receiving, at the SmartSwitch network device from a device that is remote from the virtual computing environment, a data packet addressed to an endpoint in a virtual network hosted by one of the virtual machines. In an embodiment, the data packet comprises an identifier indicative of the remote device.

(57) Operation **601** may be followed by operation **603**. Operation **603** illustrates based on the identifier, determining that the data packet is associated with the virtual network.

(58) Operation **603** may be followed by operation **605**. Operation **605** illustrates based on the determined association, mapping one of a plurality of policies to a data flow of the virtual network.

(59) Operation **605** may be followed by operation **607**. Operation **607** illustrates modifying, by the SDN appliance, the packet in accordance with the mapped policy; wherein the mapped policy is dynamically adjustable based on the data flow.

(60) Operation **607** may be followed by operation **609**. Operation **609** illustrates forwarding the modified packet to the endpoint in the virtual network.

(61) Turning now to FIG. **6B**, illustrated is another example operational procedure for processing data packets and implementing policies in a software defined network (SDN) of a virtual computing environment, by at least one SmartSwitch network device configured to disaggregate enforcement of policies of the SDN from hosts of the virtual computing environment. The hosts may be implemented on servers communicatively coupled to network interfaces of the SmartSwitch network device. The servers may host a plurality of virtual machines. The SmartSwitch network device may comprise a plurality of data processing units configured to implement functionality of the SmartSwitch network device. In an embodiment, the SmartSwitch network device may be situated at a tier-1 position in the virtual computing environment.

(62) Referring to FIG. **6B**, operation **621** illustrates receiving, by the SmartSwitch network device, a data packet addressed to an endpoint in a virtual network hosted by one of the virtual machines.

(63) Operation **621** may be followed by operation **623**. Operation **623** illustrates determining that the data packet is associated with the virtual network.

(64) Operation **623** may be followed by operation **625**. Operation **625** illustrates in response to determining that the data packet is associated with the virtual network, matching the packet to a data flow of the virtual network and one of a plurality of policies associated with the data flow.

(65) Operation **625** may be followed by operation **627**. Operation **627** illustrates modifying, by the SmartSwitch network device, the packet in accordance with the one policy matched with the packet; wherein the one policy matched with the packet is dynamically adjustable based on the data flow.

(66) Operation **627** may be followed by operation **629**. Operation **629** illustrates forwarding the modified packet to the endpoint in the virtual network.

(67) FIG. **7** illustrates an example computing environment in which the embodiments described herein may be implemented. FIG. **7** illustrates a service provider **700** that is configured to provide computing resources to users at user site **740**. The user site **740** may have user computers that may

access services provided by service provider **700** via a network **730**. The computing resources provided by the service provider **700** may include various types of resources, such as computing resources, data storage resources, data communication resources, and the like. For example, computing resources may be available as virtual machines. The virtual machines may be configured to execute applications, including Web servers, application servers, media servers, database servers, and the like. Data storage resources may include file storage devices, block storage devices, and the like. Networking resources may include virtual networking, software load balancer, and the like.

(68) Service provider **700** may have various computing resources including servers, routers, and other devices that may provide remotely accessible computing and network resources using, for example, virtual machines. Other resources that may be provided include data storage resources. Service provider **700** may also execute functions that manage and control allocation of network resources, such as a network manager **770**.

(69) Network **730** may, for example, be a publicly accessible network of linked networks and may be operated by various entities, such as the Internet. In other embodiments, network **730** may be a private network, such as a dedicated network that is wholly or partially inaccessible to the public. Network **730** may provide access to computers and other devices at the user site **740**.

(70) FIG. **8** illustrates an example computing environment in which the embodiments described herein may be implemented. FIG. **8** illustrates a data center **800** that is configured to provide computing resources to users **800a**, **800b**, or **800c** (which may be referred herein singularly as “a user **800**” or in the plural as “the users **800**”) via user computers **808a**, **808b**, and **808c** (which may be referred herein singularly as “a computer **808**” or in the plural as “the computers **808**”) via a communications network **880**. The computing resources provided by the data center **800** may include various types of resources, such as computing resources, data storage resources, data communication resources, and the like. Each type of computing resource may be general-purpose or may be available in a number of specific configurations. For example, computing resources may be available as virtual machines. The virtual machines may be configured to execute applications, including Web servers, application servers, media servers, database servers, and the like. Data storage resources may include file storage devices, block storage devices, and the like. Each type or configuration of computing resource may be available in different configurations, such as the number of processors, and size of memory and/or storage capacity. The resources may in some embodiments be offered to clients in units referred to as instances, such as virtual machine instances or storage instances. A virtual computing instance may be referred to as a virtual machine and may, for example, comprise one or more servers with a specified computational capacity (which may be specified by indicating the type and number of CPUs, the main memory size and so on) and a specified software stack (e.g., a particular version of an operating system, which may in turn run on top of a hypervisor).

(71) Data center **800** may correspond to service provider **100** in FIGS. **1** and **8**, or edge site **150** of FIG. **8**. Data center **800** may include servers **886a**, **886b**, and **886c** (which may be referred to herein singularly as “a server **886**” or in the plural as “the servers **886**”) that may be standalone or installed in server racks, and provide computing resources available as virtual machines **888a** and **888b** (which may be referred to herein singularly as “a virtual machine **888**” or in the plural as “the virtual machines **888**”). The virtual machines **888** may be configured to execute applications such as Web servers, application servers, media servers, database servers, and the like. Other resources that may be provided include data storage resources (not shown on FIG. **8**) and may include file storage devices, block storage devices, and the like. Servers **886** may also execute functions that manage and control allocation of resources in the data center, such as a controller **885**. Controller **885** may be a fabric controller or another type of program configured to manage the allocation of virtual machines on servers **886**.

(72) Referring to FIG. **8**, communications network **880** may, for example, be a publicly accessible

network of linked networks and may be operated by various entities, such as the Internet. In other embodiments, communications network **880** may be a private network, such as a corporate network that is wholly or partially inaccessible to the public.

(73) Communications network **880** may provide access to computers **808**. Computers **808** may be computers utilized by users **800**. Computer **808a**, **808b** or **808c** may be a server, a desktop or laptop personal computer, a tablet computer, a smartphone, a set-top box, or any other computing device capable of accessing data center **800**. User computer **808a** or **808b** may connect directly to the Internet (e.g., via a cable modem). User computer **808c** may be internal to the data center **800** and may connect directly to the resources in the data center **800** via internal networks. Although only three user computers **808a**, **808b**, and **808c** are depicted, it should be appreciated that there may be multiple user computers.

(74) Computers **808** may also be utilized to configure aspects of the computing resources provided by data center **800**. For example, data center **800** may provide a Web interface through which aspects of its operation may be configured through the use of a Web browser application program executing on user computer **808**. Alternatively, a stand-alone application program executing on user computer **808** may be used to access an application programming interface (API) exposed by data center **800** for performing the configuration operations.

(75) Servers **886** may be configured to provide the computing resources described above. One or more of the servers **886** may be configured to execute a manager **830a** or **830b** (which may be referred herein singularly as “a manager **830**” or in the plural as “the managers **830**”) configured to execute the virtual machines. The managers **830** may be a virtual machine monitor (VMM), fabric controller, or another type of program configured to enable the execution of virtual machines **888** on servers **886**, for example.

(76) It should be appreciated that although the embodiments disclosed above are discussed in the context of virtual machines, other types of implementations can be utilized with the concepts and technologies disclosed herein.

(77) In the example data center **800** shown in FIG. **8**, a network device **888** may be utilized to interconnect the servers **886a** and **886b**. Network device **888** may comprise one or more switches, routers, or other network devices. Network device **888** may also be connected to gateway **840**, which is connected to communications network **880**. Network device **888** may facilitate communications within networks in data center **800**, for example, by forwarding packets or other data communications as appropriate based on characteristics of such communications (e.g., header information including source and/or destination addresses, protocol identifiers, etc.) and/or the characteristics of the private network (e.g., routes based on network topology, etc.). It will be appreciated that, for the sake of simplicity, various aspects of the computing systems and other devices of this example are illustrated without showing certain conventional details. Additional computing systems and other devices may be interconnected in other embodiments and may be interconnected in different ways.

(78) It should be appreciated that the network topology illustrated in FIG. **8** has been greatly simplified and that many more networks and networking devices may be utilized to interconnect the various computing systems disclosed herein. These network topologies and devices should be apparent to those skilled in the art.

(79) It should also be appreciated that data center **800** described in FIG. **8** is merely illustrative and that other implementations might be utilized. Additionally, it should be appreciated that the functionality disclosed herein might be implemented in software, hardware or a combination of software and hardware. Other implementations should be apparent to those skilled in the art. It should also be appreciated that a server, gateway, or other computing device may comprise any combination of hardware or software that can interact and perform the described types of functionality, including without limitation desktop or other computers, database servers, network storage devices and other network devices, PDAs, tablets, smartphone, Internet appliances,

television-based systems (e.g., using set top boxes and/or personal/digital video recorders), and various other consumer products that include appropriate communication capabilities. In addition, the functionality provided by the illustrated modules may in some embodiments be combined in fewer modules or distributed in additional modules. Similarly, in some embodiments the functionality of some of the illustrated modules may not be provided and/or other additional functionality may be available.

(80) In some embodiments, aspects of the present disclosure may be implemented in a mobile edge computing (MEC) environment implemented in conjunction with a 4G, 5G, or other cellular network. MEC is a type of edge computing that uses cellular networks and 5G and enables a data center to extend cloud services to local deployments using a distributed architecture that provide federated options for local and remote data and control management. MEC architectures may be implemented at cellular base stations or other edge nodes and enable operators to host content closer to the edge of the network, delivering high-bandwidth, low-latency applications to end users. For example, the cloud provider's footprint may be co-located at a carrier site (e.g., carrier data center), allowing for the edge infrastructure and applications to run closer to the end user via the 5G network.

(81) FIG. 9 illustrates a general-purpose computing device **900**. In the illustrated embodiment, computing device **900** includes one or more processors **910a**, **910b**, and/or **910n** (which may be referred herein singularly as “a processor **910**” or in the plural as “the processors **910**”) coupled to a system memory **99** via an input/output (I/O) interface **930**. Computing device **900** further includes a network interface **940** coupled to I/O interface **930**.

(82) In various embodiments, computing device **900** may be a uniprocessor system including one processor **910** or a multiprocessor system including several processors **910** (e.g., two, four, eight, or another suitable number). Processors **910** may be any suitable processors capable of executing instructions. For example, in various embodiments, processors **910** may be general-purpose or embedded processors implementing any of a variety of instruction set architectures (ISAs), such as the x99, PowerPC, SPARC, or MIPS ISAs, or any other suitable ISA. In multiprocessor systems, each of processors **910** may commonly, but not necessarily, implement the same ISA.

(83) System memory **99** may be configured to store instructions and data accessible by processor(s) **910**. In various embodiments, system memory **99** may be implemented using any suitable memory technology, such as static random access memory (SRAM), synchronous dynamic RAM (SDRAM), nonvolatile/Flash-type memory, or any other type of memory. In the illustrated embodiment, program instructions and data implementing one or more desired functions, such as those methods, techniques and data described above, are shown stored within system memory **920** as code **925** and data **929**.

(84) In one embodiment, I/O interface **930** may be configured to coordinate I/O traffic between the processor **910**, system memory **99**, and any peripheral devices in the device, including network interface **940** or other peripheral interfaces. In some embodiments, I/O interface **930** may perform any necessary protocol, timing, or other data transformations to convert data signals from one component (e.g., system memory **920**) into a format suitable for use by another component (e.g., processor **910**). In some embodiments, I/O interface **930** may include support for devices attached through various types of peripheral buses, such as a variant of the Peripheral Component Interconnect (PCI) bus standard or the Universal Serial Bus (USB) standard, for example. In some embodiments, the function of I/O interface **930** may be split into two or more separate components. Also, in some embodiments some or all of the functionality of I/O interface **930**, such as an interface to system memory **920**, may be incorporated directly into processor **910**.

(85) Network interface **940** may be configured to allow data to be exchanged between computing device **900** and other device or devices **990** attached to a network or network(s) **990**, such as other computer systems or devices as illustrated in FIGS. 1 through 5, for example. In various embodiments, network interface **940** may support communication via any suitable wired or

wireless general data networks, such as types of Ethernet networks, for example. Additionally, network interface **940** may support communication via telecommunications/telephony networks such as analog voice networks or digital fiber communications networks, via storage area networks such as Fibre Channel SANs or via any other suitable type of network and/or protocol.

(86) In some embodiments, system memory **920** may be one embodiment of a computer-accessible medium configured to store program instructions and data as described above for the Figures for implementing embodiments of the corresponding methods and apparatus. However, in other embodiments, program instructions and/or data may be received, sent or stored upon different types of computer-accessible media. A computer-accessible medium may include non-transitory storage media or memory media, such as magnetic or optical media, e.g., disk or DVD/CD coupled to computing device **900** via I/O interface **930**. A non-transitory computer-accessible storage medium may also include any volatile or non-volatile media, such as RAM (e.g. SDRAM, DDR SDRAM, RDRAM, SRAM, etc.), ROM, etc., that may be included in some embodiments of computing device **900** as system memory **920** or another type of memory. Further, a computer-accessible medium may include transmission media or signals such as electrical, electromagnetic or digital signals, conveyed via a communication medium such as a network and/or a wireless link, such as may be implemented via network interface **940**. Portions or all of multiple computing devices, such as those illustrated in FIG. **9**, may be used to implement the described functionality in various embodiments; for example, software components running on a variety of different devices and servers may collaborate to provide the functionality. In some embodiments, portions of the described functionality may be implemented using storage devices, network devices, or special-purpose computer systems, in addition to or instead of being implemented using general-purpose computer systems. The term “computing device,” as used herein, refers to at least all these types of devices and is not limited to these types of devices.

(87) Various storage devices and their associated computer-readable media provide non-volatile storage for the computing devices described herein. Computer-readable media as discussed herein may refer to a mass storage device, such as a solid-state drive, a hard disk or CD-ROM drive. However, it should be appreciated by those skilled in the art that computer-readable media can be any available computer storage media that can be accessed by a computing device.

(88) By way of example, and not limitation, computer storage media may include volatile and non-volatile, removable and non-removable media implemented in any method or technology for storage of information such as computer-readable instructions, data structures, program modules or other data. For example, computer media includes, but is not limited to, RAM, ROM, EPROM, EEPROM, flash memory or other solid state memory technology, CD-ROM, digital versatile disks (“DVD”), HD-DVD, BLU-RAY, or other optical storage, magnetic cassettes, magnetic tape, magnetic disk storage or other magnetic storage devices, or any other medium which can be used to store the desired information and which can be accessed by the computing devices discussed herein. For purposes of the claims, the phrase “computer storage medium,” “computer-readable storage medium” and variations thereof, does not include waves, signals, and/or other transitory and/or intangible communication media, per se.

(89) Encoding the software modules presented herein also may transform the physical structure of the computer-readable media presented herein. The specific transformation of physical structure may depend on various factors, in different implementations of this description. Examples of such factors may include, but are not limited to, the technology used to implement the computer-readable media, whether the computer-readable media is characterized as primary or secondary storage, and the like. For example, if the computer-readable media is implemented as semiconductor-based memory, the software disclosed herein may be encoded on the computer-readable media by transforming the physical state of the semiconductor memory. For example, the software may transform the state of transistors, capacitors, or other discrete circuit elements constituting the semiconductor memory. The software also may transform the physical state of such

components in order to store data thereupon.

(90) As another example, the computer-readable media disclosed herein may be implemented using magnetic or optical technology. In such implementations, the software presented herein may transform the physical state of magnetic or optical media, when the software is encoded therein. These transformations may include altering the magnetic characteristics of particular locations within given magnetic media. These transformations also may include altering the physical features or characteristics of particular locations within given optical media, to change the optical characteristics of those locations. Other transformations of physical media are possible without departing from the scope and spirit of the present description, with the foregoing examples provided only to facilitate this discussion.

(91) In light of the above, it should be appreciated that many types of physical transformations take place in the disclosed computing devices in order to store and execute the software components and/or functionality presented herein. It is also contemplated that the disclosed computing devices may not include all of the illustrated components shown in FIG. 9, may include other components that are not explicitly shown in FIG. 9, or may utilize an architecture completely different than that shown in FIG. 9.

(92) Although the various configurations have been described in language specific to structural features and/or methodological acts, it is to be understood that the subject matter defined in the appended representations is not necessarily limited to the specific features or acts described. Rather, the specific features and acts are disclosed as example forms of implementing the claimed subject matter.

(93) Conditional language used herein, such as, among others, “can,” “could,” “might,” “may,” “e.g.,” and the like, unless specifically stated otherwise, or otherwise understood within the context as used, is generally intended to convey that certain embodiments include, while other embodiments do not include, certain features, elements, and/or steps. Thus, such conditional language is not generally intended to imply that features, elements, and/or steps are in any way required for one or more embodiments or that one or more embodiments necessarily include logic for deciding, with or without author input or prompting, whether these features, elements, and/or steps are included or are to be performed in any particular embodiment. The terms “comprising,” “including,” “having,” and the like are synonymous and are used inclusively, in an open-ended fashion, and do not exclude additional elements, features, acts, operations, and so forth. Also, the term “or” is used in its inclusive sense (and not in its exclusive sense) so that when used, for example, to connect a list of elements, the term “or” means one, some, or all of the elements in the list.

(94) While certain example embodiments have been described, these embodiments have been presented by way of example only, and are not intended to limit the scope of the inventions disclosed herein. Thus, nothing in the foregoing description is intended to imply that any particular feature, characteristic, step, module, or block is necessary or indispensable. Indeed, the novel methods and systems described herein may be embodied in a variety of other forms; furthermore, various omissions, substitutions and changes in the form of the methods and systems described herein may be made without departing from the spirit of the inventions disclosed herein. The accompanying claims and their equivalents are intended to cover such forms or modifications as would fall within the scope and spirit of certain of the inventions disclosed herein.

(95) It should be appreciated any reference to “first,” “second,” etc. items and/or abstract concepts within the description is not intended to and should not be construed to necessarily correspond to any reference of “first,” “second,” etc. elements of the claims. In particular, within this Summary and/or the following Detailed Description, items and/or abstract concepts such as, for example, individual computing devices and/or operational states of the computing cluster may be distinguished by numerical designations without such designations corresponding to the claims or even other paragraphs of the Summary and/or Detailed Description. For example, any designation

of a “first operational state” and “second operational state” of the computing cluster within a paragraph of this disclosure is used solely to distinguish two different operational states of the computing cluster within that specific paragraph—not any other paragraph and particularly not the claims.

(96) Although the various techniques have been described in language specific to structural features and/or methodological acts, it is to be understood that the subject matter defined in the appended representations is not necessarily limited to the specific features or acts described. Rather, the specific features and acts are disclosed as example forms of implementing the claimed subject matter.

(97) The disclosure presented herein also encompasses the subject matter set forth in the following clauses:

(98) Clause 1: A method for processing data packets and implementing policies in a software defined network (SDN) of a virtual computing environment, by at least one SmartSwitch network device configured to disaggregate enforcement of policies of the SDN from hosts of the virtual computing environment, the hosts implemented on servers communicatively coupled to network interfaces of the SmartSwitch network device, the servers hosting a plurality of virtual machines, the SmartSwitch network device comprising a plurality of data processing units configured to implement functionality of the SmartSwitch network device, the SmartSwitch network device situated at a tier-1 position in the virtual computing environment, the method comprising: receiving, by the SmartSwitch network device, a data packet addressed to an endpoint in a virtual network hosted by one of the virtual machines; determining that the data packet is associated with the virtual network; in response to determining that the data packet is associated with the virtual network, matching the packet to a data flow of the virtual network and one of a plurality of policies associated with the data flow; modifying, by the SmartSwitch network device, the packet in accordance with the one policy matched with the packet; wherein the one policy matched with the packet is dynamically adjustable based on the data flow; and forwarding the modified packet to the endpoint in the virtual network.

(99) Clause 2: The method of clause 1, wherein the SmartSwitch network device is interconnected in a Clos configuration with tier-0 network devices.

(100) Clause 3: The method of any of clauses 1-2, wherein the SmartSwitch network device is communicatively coupled to at least two top-of-rack switches that are situated in tier-0 positions in the virtual computing environment.

(101) Clause 4: The method of any of clauses 1-3, wherein the SmartSwitch network device is part of a disaggregated cluster of computing devices.

(102) Clause 5: The method of any of clauses 1-4, wherein the disaggregated cluster is a GPU cluster, bare metal cluster, or AI cluster.

(103) Clause 6: The method of any of clauses 1-5, wherein the SmartSwitch network devices are communicatively coupled to a plurality of switches at a tier-2 position in the virtual computing environment.

(104) Clause 7: A system comprising: a plurality of servers communicatively coupled to at least one SmartSwitch network device configured to disaggregate enforcement of policies of a SDN of a virtual computing environment from hosts of the virtual computing environment, the hosts implemented on servers communicatively coupled to network interfaces of the SmartSwitch network device, the servers hosting a plurality of virtual machines, the SmartSwitch network device comprising a plurality of hardware-based processing units configured to implement functionality of the SmartSwitch network device, the SmartSwitch network device situated at a tier-1 position in the virtual computing environment; the system configured to: receive a data packet addressed to an endpoint in a virtual network hosted by one of the virtual machines, the data packet comprising an identifier indicative of a source and destination of the data packet; based on the identifier: determining that the data packet is associated with the virtual network; and based on the

determining that the data packet is associated with the virtual network, matching the data packet to one of a plurality of policies associated with a data flow of the virtual network; modifying the data packet in accordance with the matched policy; wherein the matched policy is dynamically adjustable based on the data flow; and forwarding the modified packet to the endpoint in the virtual network.

(105) Clause 8: The system of clause 7, wherein the SmartSwitch network device is interconnected in a Clos configuration with tier-0 network devices.

(106) Clause 9: The system of any of clauses 7 and 8, wherein the SmartSwitch network device is communicatively coupled to at least two top-of-rack switches.

(107) Clause 10: The system of any clauses 7-9, wherein the SmartSwitch network device is a part of a disaggregated cluster.

(108) Clause 11: The hardware-based networking device of any clauses 7-10, wherein the disaggregated cluster is a GPU cluster, bare metal cluster, or AI cluster.

(109) Clause 12: The hardware-based networking device of any clauses 7-11, wherein the SmartSwitch network device is configured to apply policies of the virtual computing environment to data traffic on the virtual network after the data traffic leaves its source and before the data traffic reaches its destination.

(110) Clause 13: A network device configured to disaggregate enforcement of policies of a software defined network (SDN) of a virtual computing environment from hosts of the virtual computing environment, the hosts implemented on servers communicatively coupled to network interfaces of the network device, the servers hosting a plurality of virtual machines, the network device comprising a plurality of processing units configured to implement functionality of the network device, the network device configured to be situated at a tier-1 position in the virtual computing environment, the network device configured to: receive a data packet addressed to an endpoint in a virtual network hosted by one of the virtual machines; determining that the data packet is associated with the virtual network; and based on determining that the data packet is associated with the virtual network, matching the data packet to one of a plurality of policies associated with a data flow of the virtual network; modifying the packet in accordance with the matched policy; wherein the matched policy is dynamically adjustable based on the data flow; and forwarding the modified packet to the endpoint in the virtual network.

(111) Clause 14: The network device of clause 13, wherein the network device is interconnected in a Clos configuration with tier-0 network devices.

(112) Clause 15: The network device of any of clauses 13 and 14, wherein the network device is communicatively coupled to at least two top-of-rack switches.

(113) Clause 16: The network device of any of the clauses 13-15, wherein the network device is a part of a disaggregated cluster.

(114) Clause 17: The network device of any of the clauses 13-16, wherein the disaggregated cluster is a GPU cluster, bare metal cluster, or AI cluster.

(115) Clause 18: The network device of any of the clauses 13-17, wherein the network device is configured to apply policies of the virtual computing environment to data traffic on the virtual network after the data traffic leaves its source and before the data traffic reaches its destination.

(116) Clause 19: The network device of any of the clauses 13-18, wherein the network device is communicatively coupled to two top-of-rack switches so that each of the of-rack switches have a switchable communications path to the network device.

(117) Clause 20: The network device of any of the clauses 13-19, wherein the network device is configured with SDN agents configured to manage functionality of the network device.

Claims

1. A method for processing data packets and implementing policies in a software defined network (SDN) of a virtual computing environment, the method comprising: receiving, by a SmartSwitch network switch device, a data packet addressed to an endpoint in a virtual network hosted by one of a plurality of virtual machines of the virtual computing environment, the SmartSwitch network switch device configured to disaggregate enforcement of policies of the SDN from servers hosting the virtual machines, the servers communicatively coupled to network interfaces of the SmartSwitch network switch device, the SmartSwitch network switch device comprising a plurality of data processing units configured to implement functionality of the SmartSwitch network switch device, the SmartSwitch network switch device interconnected in a Clos configuration and situated at a tier-1 switch position in the SDN; determining that the data packet is associated with the virtual network; in response to determining that the data packet is associated with the virtual network, matching the packet to a data flow of the virtual network and one of a plurality of policies associated with the data flow; modifying, by the SmartSwitch network switch device, the packet by applying the one policy to the packet; wherein the one policy matched with the packet is dynamically adjustable based on the data flow; and forwarding the modified packet to the endpoint in the virtual network.
2. The method of claim 1, wherein the SmartSwitch network switch device is interconnected with tier-0 network devices.
3. The method of claim 1, wherein the SmartSwitch network switch device is communicatively coupled to at least two top-of-rack switches that are situated in tier-0 positions in the virtual computing environment.
4. The method of claim 1, wherein the SmartSwitch network switch device is part of a disaggregated cluster of computing devices.
5. The method of claim 4, wherein the disaggregated cluster is a GPU cluster, bare metal cluster, or AI cluster.
6. The method of claim 1, wherein the SmartSwitch network switch devices are communicatively coupled to a plurality of switches at a tier-2 position in the virtual computing environment.
7. A system comprising: a plurality of servers communicatively coupled to at least one SmartSwitch network switch device configured to disaggregate enforcement of policies of a SDN of a virtual computing environment from servers hosting a plurality of virtual machines, the SmartSwitch network switch device comprising a plurality of data processing units configured to implement functionality of the SmartSwitch network switch device, the SmartSwitch network switch device interconnected in a Clos configuration and situated at a tier-1 position in the SDN; the SmartSwitch network switch device configured to: receive a data packet addressed to an endpoint in a virtual network hosted by one of the plurality of virtual machines, the data packet comprising an identifier indicative of a source and destination of the data packet; based on the identifier: determine that the data packet is associated with the virtual network; and based on the determining that the data packet is associated with the virtual network, match the data packet to one of a plurality of policies associated with a data flow of the virtual network; modify the data packet by applying the one policy to the data packet; wherein the one policy is dynamically adjustable based on the data flow; and forward the modified packet to the endpoint in the virtual network.
8. The system of claim 7, wherein the SmartSwitch network switch device is interconnected with tier-0 network devices.
9. The system of claim 7, wherein the SmartSwitch network switch device is communicatively coupled to at least two top-of-rack switches.
10. The system of claim 7, wherein the SmartSwitch network switch device is a part of a disaggregated cluster.
11. The system of claim 10, wherein the disaggregated cluster is a GPU cluster, bare metal cluster, or AI cluster.

12. The system of claim 7, wherein the SmartSwitch network switch device is configured to apply policies of the virtual computing environment to data traffic on the virtual network after the data traffic leaves its source and before the data traffic reaches its destination.
13. A SmartSwitch network switch device configured to disaggregate enforcement of policies of a software defined network (SDN) of a virtual computing environment from servers communicatively coupled to network interfaces of the SmartSwitch network switch device, the servers hosting a plurality of virtual machines, the SmartSwitch network switch device comprising a plurality of data processing units configured to implement functionality of the SmartSwitch network switch device, the SmartSwitch network switch device configured to be interconnected in a Clos configuration and situated at a tier-1 position in the SDN, the plurality of data processing units configuring the SmartSwitch network switch device to perform operations comprising: receiving a data packet addressed to an endpoint in a virtual network hosted by one of the plurality of virtual machines; determining that the data packet is associated with the virtual network; and based on determining that the data packet is associated with the virtual network, matching the data packet to one of a plurality of policies associated with a data flow of the virtual network; modifying the packet by applying the one policy to the data packet; wherein the one policy is dynamically adjustable based on the data flow; and forwarding the modified packet to the endpoint in the virtual network.
14. The SmartSwitch network switch device of claim 13, wherein the SmartSwitch network switch device is interconnected with tier-0 network devices.
15. The SmartSwitch network switch device of claim 13, wherein the SmartSwitch network switch device is communicatively coupled to at least two top-of-rack switches.
16. The SmartSwitch network switch device of claim 13, wherein the SmartSwitch network switch device is a part of a disaggregated cluster.
17. The SmartSwitch network switch device of claim 16, wherein the disaggregated cluster is a GPU cluster, bare metal cluster, or AI cluster.
18. The SmartSwitch network switch device of claim 13, wherein the SmartSwitch network switch device is configured to apply policies of the virtual computing environment to data traffic on the virtual network after the data traffic leaves its source and before the data traffic reaches its destination.
19. The SmartSwitch network switch device of claim 16, wherein the SmartSwitch network switch device is communicatively coupled to two top-of-rack switches so that each of the of-rack switches have a switchable communications path to the SmartSwitch network switch device.
20. The SmartSwitch network switch device of claim 13, wherein the SmartSwitch network switch device is configured with SDN agents configured to manage functionality of the SmartSwitch network switch device.
-