

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12395448
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Hunt; Gwendolyn Meredith et al.

Agent message delivery fairness

Abstract

Apparatus and methods are disclosed for generating, sending, and receiving messages in a networked environment using autonomous (or semi-autonomous) agents. In one example of the disclosed technology, a method of controlling message flow in a computer network comprising a plurality of agents, agent data consumers, and an agent message bridge configured to send messages by receiving a set of messages, at least some of the messages including a message type, queuing the set of messages in a spooler that includes an indication of the respective message type for each of the messages, receive an indication that sending some of the messages queued in the spooler should be delayed for one or more indicated message types, and sending at least one of the messages to a selected one or more of the agent data consumers, the sent messages not being of the indicated message types.

Inventors: Hunt; Gwendolyn Meredith (Portland, OR), Jansen; Benjamin J. (Tigard, OR), Delaney; Eric L. (Portland, OR), Larson; Ryan (Keizer, OR), Thomas; Mitch (Portland, OR), McFeely; Brian (Orinda, CA)

Applicant: Tripwire, Inc. (Portland, OR)

Family ID: 1000008763289

Assignee: TRIPWIRE, INC. (Portland, OR)

Appl. No.: 18/542042

Filed: December 15, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240113984 A1	Apr. 04, 2024

Related U.S. Application Data

continuation-in-part parent-doc US 14303431 20140612 US 9634951 20170425 child-doc US 14818169
division parent-doc US 17497745 20211008 US 11863460 child-doc US 18542042
division parent-doc US 16414390 20190516 US 11159439 20211026 child-doc US 17497745
division parent-doc US 14818169 20150804 US 10313257 20190604 child-doc US 16414390
us-provisional-application US 62032975 20140804

Publication Classification

Int. Cl.: **H04L47/50** (20220101); **G06F11/14** (20060101); **H04L9/32** (20060101); **H04L9/40** (20220101); **H04L47/34** (20220101)

U.S. Cl.:

CPC **H04L47/50** (20130101); **G06F11/1464** (20130101); **H04L9/3263** (20130101); **H04L47/34** (20130101); **H04L63/0428** (20130101); G06F2201/84 (20130101)

Field of Classification Search

CPC: H04L (47/50); H04L (9/3263); H04L (47/34); H04L (63/0428); G06F (11/1464); G06F (2201/84)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6092109	12/1999	Goto	709/224	H04L 1/1835
6314190	12/2000	Zimmermann	380/282	H04L 9/0894
6785821	12/2003	Teal	N/A	N/A
6941358	12/2004	Joiner et al.	N/A	N/A
7243348	12/2006	Good et al.	N/A	N/A
7246156	12/2006	Ginter	709/227	H04L 63/1441
7316016	12/2007	Difalco	N/A	N/A
7360099	12/2007	Difalco et al.	N/A	N/A
7587754	12/2008	Difalco et al.	N/A	N/A
7620715	12/2008	Difalco et al.	N/A	N/A
7765460	12/2009	Difalco et al.	N/A	N/A
7822724	12/2009	Difalco et al.	N/A	N/A
7979245	12/2010	Bourlatchkov	703/2	G06F 11/3409
8079081	12/2010	Lavrik et al.	N/A	N/A
8140635	12/2011	Difalco	N/A	N/A
8176158	12/2011	Difalco et al.	N/A	N/A
8346935	12/2012	Mayo et al.	N/A	N/A
9634951	12/2016	Hunt et al.	N/A	N/A
10313257	12/2018	Hunt et al.	N/A	N/A
10764257	12/2019	Hunt et al.	N/A	N/A
11159439	12/2020	Hunt et al.	N/A	N/A
2002/0196330	12/2001	Park et al.	N/A	N/A

2004/0024843	12/2003	Smith	N/A	N/A
2004/0075738	12/2003	Burke et al.	N/A	N/A
2004/0083373	12/2003	Perkins et al.	N/A	N/A
2004/0122962	12/2003	Difalco et al.	N/A	N/A
2004/0139166	12/2003	Collison	N/A	N/A
2004/0151114	12/2003	Ruutu	N/A	N/A
2004/0230903	12/2003	Elza et al.	N/A	N/A
2005/0210331	12/2004	Connelly et al.	N/A	N/A
2006/0206883	12/2005	Sabbouh	N/A	N/A
2006/0242277	12/2005	Torrence et al.	N/A	N/A
2007/0005665	12/2006	Vaitzblit et al.	N/A	N/A
2007/0016910	12/2006	Roberts	N/A	N/A
2007/0043786	12/2006	Difalco et al.	N/A	N/A
2007/0061815	12/2006	Graham et al.	N/A	N/A
2007/0124255	12/2006	Difalco et al.	N/A	N/A
2007/0239862	12/2006	Bronez et al.	N/A	N/A
2008/0016501	12/2007	Muhlestein et al.	N/A	N/A
2008/0021912	12/2007	Seligman et al.	N/A	N/A
2008/0168420	12/2007	Sabbouh	N/A	N/A
2009/0193298	12/2008	Mukherjee	N/A	N/A
2009/0228577	12/2008	Webb-Johnson	N/A	N/A
2010/0005107	12/2009	Difalco	N/A	N/A
2010/0043066	12/2009	Miliefsky	N/A	N/A
2011/0066951	12/2010	Ward-Karet et al.	N/A	N/A
2011/0099236	12/2010	Vuong et al.	N/A	N/A
2011/0197094	12/2010	Wagner	N/A	N/A
2011/0197189	12/2010	Wagner et al.	N/A	N/A
2011/0197205	12/2010	Wagner et al.	N/A	N/A
2011/0208841	12/2010	Robertson et al.	N/A	N/A
2011/0258268	12/2010	Banks et al.	N/A	N/A
2012/0005688	12/2011	Kesavan et al.	N/A	N/A
2012/0023076	12/2011	Torrence et al.	N/A	N/A
2012/0089568	12/2011	Manley et al.	N/A	N/A
2012/0179805	12/2011	Difalco	N/A	N/A
2012/0210434	12/2011	Curtis et al.	N/A	N/A
2012/0259912	12/2011	Kruse et al.	N/A	N/A
2014/0237145	12/2013	Mints	N/A	N/A
2015/0293955	12/2014	Dickey	N/A	N/A

OTHER PUBLICATIONS

ITU-T Recommendation X.667, “Information technology—Open Systems Interconnection-Procedures for the operation of OSI Registration Authorities: Generation and registration of Universally Unique Identifiers (UUIDs) and their use as ASN. 1 object identifier components,” (document marked Sep. 2004), 34 pages. cited by applicant

U.S. Appl. No. 17/001,402, unpublished application, Hunt et al. cited by applicant

Primary Examiner: Ndiaye; Cheikh T

Attorney, Agent or Firm: Fredrikson & Byron, P.A.

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a divisional of U.S. patent application Ser. No. 17/497,745, filed on Oct. 8, 2021, which is a divisional of U.S. patent application Ser. No. 16/414,390, filed on May 16, 2019, which is a divisional of U.S. patent application Ser. No. 14/818,169, filed on Aug. 4, 2015 (now U.S. Pat. No. 10,313,257), which is a continuation-in-part of U.S. patent application Ser. No. 14/303,431, entitled “Autonomous Agent Messaging,” filed on Jun. 12, 2014 (now U.S. Pat. No. 9,634,951). This application U.S. patent application Ser. No. 14/818,169 also claims the benefit of U.S. Provisional Patent Application No. 62/032,975, entitled “Agent Message Delivery Fairness,” filed on Aug. 4, 2014. The entire disclosures of U.S. patent application Ser. Nos. 16/414,390, 14/818,169, 14/303,431, and 62/032,975 are incorporated fully by reference herein for all purposes.

BACKGROUND

(1) Existing compliance and security systems for enterprise environments employ the use of agents for monitoring system integrity and reporting changes to one or more centralized compliance servers. However, existing agents are limited in a number of ways. Known examples of compliance agents are implemented as a monolithic Java agent that must be connected to a server almost continuously. Further, existing agents do not allow for plugins, quarantine, prioritizing of messaging, or disconnected operation. Accordingly, there is ample opportunity for improvement in the implementation of agents for monitoring system integrity and reporting changes.

SUMMARY

(2) Apparatus and methods are disclosed for generating, sending, and receiving messages in a networked environment using autonomous (or semi-autonomous) agents. Techniques disclosed herein for sending and receiving messages with agents include flow control mechanisms that allow for spooling of collected data by the agents to improve reliability in the event of network delay and outages. In some examples of the disclosed technology, agent capabilities can be readily expanded through the use of on-demand agent plugins for facilitating data collection and which further allow separation of message capabilities from particular plugin identification. In some examples of the disclosed technology, message and agent identifiers are applied to increase data reliability and to allow re-transmission of lost messages upon request.

(3) The disclosed agent platforms are designed to address one or more challenges presented in enterprise deployments of agents by, for example: reducing agent footprint, improving scalability, dealing with erratic computer networks, providing semi-autonomous; operation, and/or providing a self-correcting framework.

(4) In some examples of the disclosed technology, a method of controlling message flow in a computer network, the network comprising a plurality of agents, a plurality of agent data consumers, and an agent message bridge configured to send messages between the agents and the agent data consumers, includes with the agent bridge, receiving a set of messages, at least some of the messages including a message type, queuing the set of messages in a spooler, wherein the spooler includes an indication of the respective message type for each of the at least some of the messages, receiving an indication that sending of one or more, but not all, of the messages queued in the spooler should be delayed for one or more indicated message types, and sending at least one of the messages to a selected one or more of the agent data consumers, the sent messages not being of the indicated message types.

(5) In some examples, the method of controlling message flow further includes, with the agent bridge, monitoring a plurality of message topics or queues, each of the message topics or queues having a distinct type, each of the message topics or queues being configured to temporarily store messages received from the agents, and when the number of messages queued in a first message

topic or queue of the plurality of message topics or queues exceeds a predefined stop sending level, sending an advisory message from the agent bridge to the agents indicating that messages of the corresponding type should not be sent, and after sending the stop sending advisory message, when the number of messages queued in the first message topic or queue reaches a predefined restart sending level, sending an advisory message from the message bridge to the agents indicating that sending of messages of the corresponding type can be resumed.

(6) In some examples of the disclosed technology, a method of collecting data from an agent executing on a host computer includes collecting host data, the collecting occurring whether or not the agent can currently send data via the network connection where when the agent cannot send the data via the network connection, storing at least a portion of the collected host data in a spooler, and when the agent can send the data via the network connection, sending at least a portion of the spooled host data to at least one of the agent data consumers. In some examples, the method further includes receiving a message from an agent bridge indicating a message type to send to at least one of the agent data consumers and initiating sending of data for the indicated message type.

(7) In some examples of the disclosed technology, a system includes one or more agent data consumers, one or more agent platform servers, each of the agent platform servers including an agent bridge, the agent bridge being configured to send an advisory message to adjust the rate at which messages are sent by one or more agents, one or more agents, the agents being configured to execute on computing hosts separate from the agent platform servers, the agents being further configured to, responsive to receiving a message from the agent bridge, adjust the rate at which messages are sent to the agent bridge. In some examples of the system, one or more of the agent data consumers, agent platform servers, and/or agent bridges are implemented with computer hardware including a real or virtualized processor and memory. Each of the processors can be coupled to a computer network to enable communication of messages and other data between the consumer(s), server(s), and bridge(s).

(8) In some examples of the disclosed technology, one or more computer-readable storage media storing computer-readable instructions that when executed by a processor, cause the processor to perform any of the methods disclosed herein.

(9) This summary is provided to introduce a selection of concepts in a simplified form that are further described below in the detailed description. This summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used to limit the scope of the claimed subject matter. The foregoing and other objects, features, and advantages of the disclosed subject matter will become more apparent from the following detailed description, which proceeds with reference to the accompanying figures. Further, any trademarks used herein are the property of their respective owners.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1 illustrates an exemplary computing environment including an agent, a bridge, and a compliance server, as can be used in some embodiments of the disclosed technology.

(2) FIG. 2 is a block diagram illustrating an agent, as can be used in some embodiments of the disclosed technology.

(3) FIG. 3 is a block diagram further illustrating an agent, as can be used in some embodiments of the disclosed technology.

(4) FIG. 4 is a diagram that outlines an example of communication between an agent and an agent platform server, as can be used in some embodiments of the disclosed technology.

(5) FIG. 5 is a diagram that illustrates an example handshake, as can be used in some embodiments of the disclosed technology.

- (6) FIG. 6 is a diagram illustrating certificates that can be used in some embodiments of the disclosed technology.
- (7) FIG. 7 is a diagram that illustrates message flow between agent and bridge components, as can be used in some embodiments of the disclosed technology.
- (8) FIG. 8 is a diagram illustrating aspects of a spooler, as can be used in some embodiments of the disclosed technology.
- (9) FIG. 9 is a flow chart illustrating an example method of spooling host data, as can be performed in some examples of the disclosed technology.
- (10) FIG. 10 is a flow chart illustrating an example method of resending messages based on agent-generated sequence numbers, as can be performed in some examples of the disclosed technology.
- (11) FIG. 11 is a flow chart illustrating an example method of using a unique identifier for sending agent messages, as can be performed in some examples of the disclosed technology.
- (12) FIG. 12 is a diagram that illustrates an exemplary set of message flows, as can be performed in some examples of the disclosed technology.
- (13) FIG. 13 is a diagram that illustrates an exemplary set of message flows, as can be performed in some examples of the disclosed technology.
- (14) FIG. 14 illustrates a sequence diagram depicting an exemplary set of message flows, as can be performed in some examples of the disclosed technology.
- (15) FIGS. 15A and 15B illustrate a sequence diagram depicting an exemplary set of message flows, as can be performed in some examples of the disclosed technology.
- (16) FIGS. 16A and 16B illustrate a sequence diagram depicting an exemplary set of message flows, as can be performed in some examples of the disclosed technology.
- (17) FIG. 17 illustrates an exemplary computing environment including an agent, a bridge, and a compliance server, as can be used in some embodiments of the disclosed technology.
- (18) FIG. 18 illustrates a sequence diagram depicting an exemplary set of message flows, as can be performed in some examples of the disclosed technology.
- (19) FIG. 19 is an example description of an affinity service map, as can be used in certain examples of the disclosed technology.
- (20) FIGS. 20A and 20B illustrate an example description of an affinity service map, as can be used in certain examples of the disclosed technology.
- (21) FIG. 21 is a flow chart illustrating an example method of delaying certain messages, as can be performed in certain examples of the disclosed technology.
- (22) FIG. 22 is a block diagram illustrating a suitable computing environment for implemented some embodiments of the disclosed technology.

DETAILED DESCRIPTION

(23) I. General Considerations

- (24) This disclosure is set forth in the context of representative embodiments that are not intended to be limiting in any way.
- (25) As used in this application the singular forms “a,” “an,” and “the” include the plural forms unless the context clearly dictates otherwise. Additionally, the term “includes” means “comprises.” Further, the term “coupled” encompasses mechanical, electrical, magnetic, optical, as well as other practical ways of coupling or linking items together, and does not exclude the presence of intermediate elements between the coupled items. Furthermore, as used herein, the term “and/or” means any one item or combination of items in the phrase.
- (26) The systems, methods, and apparatus described herein should not be construed as being limiting in any way. Instead, this disclosure is directed toward all novel and non-obvious features and aspects of the various disclosed embodiments, alone and in various combinations and subcombinations with one another. The disclosed systems, methods, and apparatus are not limited to any specific aspect or feature or combinations thereof, nor do the disclosed things and methods require that any one or more specific advantages be present or problems be solved. Furthermore,

any features or aspects of the disclosed embodiments can be used in various combinations and subcombinations with one another.

(27) Although the operations of some of the disclosed methods are described in a particular, sequential order for convenient presentation, it should be understood that this manner of description encompasses rearrangement, unless a particular ordering is required by specific language set forth below. For example, operations described sequentially may in some cases be rearranged or performed concurrently. Moreover, for the sake of simplicity, the attached figures may not show the various ways in which the disclosed things and methods can be used in conjunction with other things and methods. Additionally, the description sometimes uses terms like “produce,” “generate,” “display,” “receive,” “evaluate,” “vulnerability,” “weakness,” “scan,” and “perform” to describe the disclosed methods. These terms are high-level abstractions of the actual operations that are performed. The actual operations that correspond to these terms will vary depending on the particular implementation and are readily discernible by one of ordinary skill in the art.

(28) Theories of operation, scientific principles, or other theoretical descriptions presented herein in reference to the apparatus or methods of this disclosure have been provided for the purposes of better understanding and are not intended to be limiting in scope. The apparatus and methods in the appended claims are not limited to those apparatus and methods that function in the manner described by such theories of operation.

(29) Any of the disclosed methods can be implemented as computer-executable instructions stored on one or more computer-readable media (e.g., non-transitory computer-readable storage media, such as one or more optical media discs, volatile memory components (such as DRAM or SRAM), or nonvolatile memory components (such as hard drives and solid state drives (SSDs))) and executed on a computer (e.g., any commercially available computer, including smart phones or other mobile devices that include computing hardware). Any of the computer-executable instructions for implementing the disclosed techniques, as well as any data created and used during implementation of the disclosed embodiments, can be stored on one or more computer-readable media (e.g., non-transitory computer-readable storage media). The computer-executable instructions can be part of, for example, a dedicated software application, or a software application that is accessed or downloaded via a web browser or other software application (such as a remote computing application). Such software can be executed, for example, on a single local computer (e.g., as an agent executing on any suitable commercially available computer) or in a network environment (e.g., via the Internet, a wide-area network, a local-area network, a client-server network (such as a cloud computing network), or other such network) using one or more network computers.

(30) For clarity, only certain selected aspects of the software-based implementations are described. Other details that are well known in the art are omitted. For example, it should be understood that the disclosed technology is not limited to any specific computer language or program. For instance, the disclosed technology can be implemented by software written in C, C++, Java, or any other suitable programming language. Likewise, the disclosed technology is not limited to any particular computer or type of hardware. Certain details of suitable computers and hardware are well-known and need not be set forth in detail in this disclosure.

(31) Furthermore, any of the software-based embodiments (comprising, for example, computer-executable instructions for causing a computer to perform any of the disclosed methods) can be uploaded, downloaded, or remotely accessed through a suitable communication means. Such suitable communication means include, for example, the Internet, the World Wide Web, an intranet, software applications, cable (including fiber optic cable), magnetic communications, electromagnetic communications (including RF, microwave, and infrared communications), electronic communications, or other such communication means.

(32) II. Exemplary Computing Network Environment

(33) FIG. 1 illustrates an exemplary computing environment **100** in which some examples of the

disclosed technology can be implemented. A number of agents **110**, **111**, and **112** are illustrated in FIG. **1**. One of the agents **110** is further detailed as shown, and includes a local agent process **120** that can manage and communicate with a number of plugins **130-135** (e.g., a file integrity monitoring (FIM) plugin **130**, a command output capture rule (COCR) plugin **131**, an Open Vulnerability Assessment Language (OVAL) plugin **132**, a Windows event log (WEL) plugin **133**, a Registry plugin **134**, and a support plugin **135**) that are configured to extend the functionality of the agent. Further details and examples of agents are discussed further below. As will be readily understood to one of ordinary skill in the relevant art, the agent technology disclosed in this paragraph is not limited to the functionality of agent plugins **130-135**, but can be adapted to specific deployments by adding other plugins or removing the depicted plugins.

(34) Each of the agents **110-112** communicates with the rest of the system depicted in the computing environment **100** via an agent platform server **150**. As shown, the agent platform server **150** includes an agent bridge **160** for sending messages to and from agents (e.g., agents **110-112**). The agent bridge **160** can send messages over a computer network to agents executing on other computers, using inter-process and/or inter-thread communication to agents executing on the same computer as the communication bridge, or by using other suitable communication means. The illustrated agent platform server **150** also includes a message broker **170** with multiple message queues **175-178** for temporarily storing messages received from and sent to, for example, the agent bridge **160**, an agent manager **180**, an affinity service **185**, and agent data consumers **190**. In some examples, the message broker **170** has a single message queue **175**. The agent platform server **150** coordinates operation of the agents by sending and receiving messages using the message broker **170**.

(35) Some agent platform server implementations can contain more than one message broker **170** organized as a network of message brokers. Additionally, some implementations can include additional instances of the agent bridge **160** or the agent manager **180**. Various combinations of message brokers, agent bridges, and agent managers can be used to support high-availability and redundant capabilities.

(36) The exemplary computing environment **100** includes a number of agent data consumers **190**, including, but not limited to, a compliance server **191**, a log server **192**, a policy server **193**, a change management server **194**, and a file integrity monitoring server **195**, an agent reconciliation server **196**, an agent provisioning server **197**, and an agent management server **198**. In some examples, different combinations of agent data consumers **190** can be deployed in the environment **100** according to the desired compliance and security applications to be performed. These combinations are not limited to a single machine. The agent bridge **160**, message broker **170**, agent manager **180**, or any combination of the agent data consumers can execute on separate computers, or separate virtual machines on a single or multiple computers. For example, the compliance server **191** can host a Compliance and Configuration Control (CCC) tool used to detect, analyze, and report on change activity in an IT infrastructure. The CCC tool can assess or receive configurations of the one or more nodes at one or more locations and determine whether the nodes comply with internal and external policies (e.g., government, regulatory, or third-party standards, such as Sarbanes-Oxley, HIPAA, ISO 27001, NIST 800, NERC, PCI, PCI-DSS, Basel II, Bill 198, CIS, DISA, FDCC, FFIEC, GCSx, GLBA, GPG 13, IBTRM, or other IT infrastructure compliance standards). The CCC tool can identify and validate changes to ensure these configurations remain in known and trusted states.

(37) In particular implementations, the CCC tool operates by capturing a baseline of server file systems, desktop file system, directory servers, databases, virtual systems, middleware applications, and/or network device configurations in a known good state. Ongoing integrity checks then compare the current states against these baselines to detect changes. The CCC tool collects information used to reconcile changes detected by the agents **110-112**, ensuring they are authorized and intended changes. The CCC tool can crosscheck detected changes with defined IT

compliance policies (e.g., using policy-based filtering), with documented change tickets in a change control management (“CCM”) system, with a list of approved changes, with automatically generated lists created by patch management and software provisioning tools, and/or against other desired and approved changes. This allows the CCC tool to automatically recognize desired changes and expose undesired changes.

(38) The CCC tool can also generate one or more reports concerning the monitored nodes showing a wide variety of information (e.g., compliance information, configuration information, usage information, etc.) The compliance-related reports generated by the CCC tool can, in some instances, comprise a score for a node that indicates the relative compliance status of the node as a numerical value in a range of possible values (e.g., a score of 1 to 100 or other such numeric or alphabetical range). The CCC tool can also apply a set of one or more tests to the nodes to evaluate the compliance status of one or more nodes. In such embodiments, the compliance-related reports generated by the CCC tool can include the number of devices that passed a particular test as well as the number of devices that failed the test. Further, the CCC tool can store detected change event data in an event log or transmit the event data as soon as it is detected or shortly after it is detected. Event logs typically comprise a list of activities and configuration changes at nodes of the IT network.

(39) An exemplary CCC tool that is suitable for use with the disclosed technology is the Tripwire® Enterprise tool available from Tripwire, Inc. The examples described below are sometimes shown or discussed as being used in connection with the Tripwire Enterprise tool. This particular usage should not be construed as limiting, however, as the disclosed technology can be adapted by those skilled in the art to help monitor and manage IT nodes using other compliance and configuration control tools as well.

(40) The compliance server **191** can also include a security information and event management (SIEM) tool that is used to centralize the storage and interpretation of events, logs, or compliance reports observed and generated in an IT management infrastructure. The event, log, and compliance report information is typically produced by other software running in the IT network. For example, CCC tools generate events that are typically kept in event logs or stored in compliance reports, as discussed above. The SIEM can be used to provide a consistent central interface that an IT administrator can use to more efficiently monitor and manage activity and configuration changes in an IT network. As needed, the IT administrator can access and use the CCC tool, which may provide deeper information than that provided by the SIEM. A SIEM tool can also integrate with external remediation, ticketing, and/or workflow tools to assist with the process of incident resolution. Furthermore, certain SIEMs include functionality for generating reports that help satisfy regulatory requirements (e.g., Sarbanes-Oxley, PCI-DSS, GLBA, or any other such requirement or standard such as any of those listed above). For these reasons, SIEM tools are becoming more widely adopted by IT administrators who desire to use a single, centralized interface for monitoring and managing their increasingly complex IT infrastructures.

(41) Logging tools can operate similarly to SIEM tools. Accordingly, for any of the embodiments disclosed below, a logging tool may take the place of a SIEM tool. For ease of readability, however, reference will typically be made to just a SIEM tool. An exemplary tool for logging and SIEM that is suitable for use with the disclosed technology is the Tripwire® Log Center tool available from Tripwire, Inc.

(42) III. Example Agent Implementation

(43) FIG. 2 is a block diagram **200** further detailing the exemplary agent **110** introduced above regarding FIG. 1. As shown in FIG. 2, the agent **110** includes one or more local agent processes **120** that interact with a number of different components (e.g., components **220**, **225**, **230**, **235**, **240**, **250**, **260**, and **270**) to perform various agent functionalities. It should be readily understood to one of ordinary skill in the art that other examples of agents can include or omit some of the components illustrated in FIG. 2.

(44) In some examples of the disclosed technology, the agent **110** provides a common platform for executing pluggable platform and/or native code in a manner that does not require a concurrently active connection to either the agent bridge **160** or agent data consumers **190**. By allowing unconnected operation, the agent **110** is better able to tolerate intermittent network connections, delays, and/or errors in the agent platform server **150**, agent data consumers **190**, or interconnecting networks.

(45) The agent **110** includes functionality for automatically adjusting the rate at which data on the host system is acquired based on, for example, currently-available host system resources including cache resources, host system workload, or other host system resources. In some examples, cached data can be resequenced based on priority changes and observed behavior of the host system. In some examples, the agent can automatically adjust and prioritize transmission of cached data to the agent bridge **160**, based on, for example, the amount of time the agent has been connected to the network, a network reconnection event, and/or using a pseudorandom number to determine when to send cached data to the agent bridge. In some examples, the adjusted rate is based on the amount of lag between messages in a spool (e.g., spooler lag can be defined by an agent as the amount of time between the oldest and newest unsent messages in a spool). In some examples, certain messages can be prioritized over others (e.g., messages carrying Security Content Automation Protocol (SCAP) data can be prioritized so that they are sent with higher priority than other types of messages).

(46) In some examples of the disclosed technology, the agent **110** is implemented in a microkernel-based operating system platform, while in other examples, the agent is implemented using a more traditional monolithic kernel. The agent can include an embedded scheduler (e.g., executed by the local agent process **120** or another process) that determines when to execute agent tasks, even when the agent is not connected to a bridge or server.

(47) In some examples, the agent **110** is a container-based agent that implements Federal Information Processing Standard (FIPS) cryptographic services for communicating and/or storing data. In some examples, information regarding FIPS containers, names, or other relevant FIPS fields are removed from data (e.g., before transmitting or storing FIPS data) to increase the difficulty of unauthorized decryption of FIPS communications and stored data.

(48) In some examples, the agent **110** includes autonomous configuration capabilities. For example, the agent **110** can determine software versions and installed hardware associated with its host system or with installed plugins and based on the determined software and hardware, negotiate a more detailed configuration with any of the agent data consumers **190**.

(49) In some examples, the agent **110** includes support for on-demand push down of plugin modules. In some examples, the agent **110** includes the capability to automatically switch to different pre-designated endpoints by automatically switching to particular ports and/or bridges.

(50) In some examples, the compliance server **191** communicates a desired spool depth to agents, which in turn adjust the rate at which data is sent to server. In some examples, when a spool associated with an agent becomes completely full, the agent can insert a mark in the spool and then, once space in the spool becomes available, peel off logs when data transmission resumes.

(51) As shown in FIG. 2, the agent **110** includes an asynchronous service module **220** for controlling and coordinating asynchronous services, for example, processing of asynchronous messages received from and sent to the agent bridge. The asynchronous service module **220** can employ a number of asynchronous input/output (I/O) threads **255** for performing these tasks.

(52) An agent information module **225** is used to send messages with information about the agent and its associated plugins, including identification information (e.g., one or more UUIDs), catalogs of available messages the agent is capable of consuming or producing, and other agent information.

(53) A message dispatcher **230** sends messages between an agent bridge (e.g., via a bridge connector) and agent plugins. In some examples, the message dispatcher **230** can send commands to an agent spooler. A message builder **235** is used to build messages sent by the message

dispatcher, including envelopes for such messages.

(54) A plugin manager **240** including a number of plugin connectors **245-247** for connecting the agent to its plugins. A thread manager **250** is used to manage agent threads (e.g., bridge writer threads, plugin manager threads, asynchronous I/O threads, or other agent threads).

(55) A bridge connector **260** is used to connect to one or more agent bridges and send messages from, for example, the message builder.

(56) A multi-file spooler **270** includes multiple spool files **275-277** that can store data from the plugin manager before the data is sent to, for example, one or more of the agent bridges.

(57) In some examples of the disclosed technology, agents are designed to provide multi-platform functionality, thus allowing developers to develop agents for, e.g., both Windows and Posix platforms concurrently.

(58) In some examples, agents and their corresponding plugins are written in C++ using multi-platform libraries and coding methodologies. In some examples, using languages such as C++ allows for a smaller agent memory footprint than agents implemented using other languages, e.g., Java.

(59) In some examples, one or more agents (e.g., agents **110-112**), agent bridges (e.g., agent bridge **160**), and/or agent data consumers **190** (e.g., compliance server **191**) can be co-located on the same computer system. In other examples, each of the agents, agent bridges, and compliance servers are installed on separate computing systems that are connected using a network or other communication means, or are installed within separate virtual machines connected on a single computing system.

(60) In some examples of the disclosed technology, the agent is executed as a non-root/non-administrator user. This provides additional security by restricting access, but in some deployments, it may be desirable to allow limited administrator access to the agent and/or a subset of agent plugins to, for example, allow access to administrator resources (e.g., to access the Windows Event Log (WEL)).

(61) The agents can communicate to the bridge using, for example, a proxy provided that supports the SOCKS5 protocol, although other protocols can be employed. In some examples, it is desirable to utilize authentication features provided by the network protocol to limit access to, for example, the bridge and/or compliance server to authenticated agents. In some examples, the SOCKS5 proxy used can be previously installed by a system administrator, and be used to support other communications unrelated to agent traffic. One desirable aspect of not including a proxy server within an agent is that the attack surface of the agent is reduced, as there is no open SOCKS5 port for attackers to attempt to attack.

(62) In some examples, the spooler **270** is supplemented by a parallel Last-In First-Out buffer (LIFO) for certain types of messages. For example, because consumers of SCAP information often prioritize the most recent data available over older data, the agent can use a LIFO as a second spool for data coming from, e.g., an OVAL plugin, such that the newest messages are transmitted to the server first.

(63) FIG. 3 is a block diagram **300** that further illustrates variations and details regarding the architecture of the exemplary agent **110** discussed above regarding FIGS. 1 and 2.

(64) A. Agent Identification

(65) In some examples of the disclosed technology, agents can use a unique identifier (e.g., a UUID (Universally Unique Identifier)), to identify themselves. The agent self-generates its unique identifier. The unique identifier is used to identify messages arriving to the bridge, as well as allowing for the routing of messages from server-side components to the agent. The unique identifier is independent of any network addresses (e.g., IPv4 or IPv6 addresses or other network addresses). In some examples, the unique identifier is associated with a set of MAC addresses associated with network interfaces discovered on the agent's host system.

(66) When one or more network addresses (e.g., an IP address) on a system change, the agent can

make note of this fact, but this does not substantially change operation of the agent. The agent can send IP addresses and associated names to the server for informational purposes, but identification of the agent by the server is primarily, if not exclusively, based on the unique identifier. An agent's identifier is not changed when the IP addresses of the agent's host system changes.

(67) B. Local Agent Process

(68) Each agent is controlled by a local agent process (e.g., local agent process **120**). The local agent process can control, for example, the bridge connector **260**, message dispatcher **230**, plugin manager **240**, and spooler **270** shown in FIG. 3.

(69) The message dispatcher **230** controls communication flow between the plugin manager **240**, the bridge connector **260**, and the spooler **270**. The message dispatcher **230** communicates with a message builder **235** for forming the message, and can include a number of bridge handlers **236** and plugin handlers **237** that are configured to operate with particular bridges and plugins, respectively.

(70) The bridge connector **260** is used to connect to and send messages to and from the agent bridge **160**. As shown in FIG. 3, the bridge connector **260** includes an asynchronous heartbeat timer **265**, and an asynchronous connection timer **266**. The asynchronous heartbeat timer **265** is used to determine intervals between sending "heartbeat" messages to the agent bridge **160**, as will be further detailed below. The asynchronous connection timer **266** is used to reset a connection attempt to the bridge if the connection is not achieved by the end of a specified timeout period.

(71) Also shown is an agent configuration manager **310**, which is responsible for reading configuration data (e.g., configuration data stored in one or more configuration files **315**) to determine how to configure the agent. Parameters that can be configured with the configuration data include, but are not limited to, timer and heartbeat time intervals, spooler configuration, plugin configuration, and security and encryption parameters. In some examples, the agent configuration manager **310** is responsible for searching for and invoking plugins by sending commands of enabled and/or disabled plugins to the plugin manager **240**.

(72) C. Agent Message Catalog

(73) The agent **110** can also be configured to publish a catalog of messages that it consumes and produces. The agent **110** does not need to publish the plugins that are consuming and/or producing information. In this way, the actual plugins being used is not shown to the consumers. If a plugin becomes disabled, then the messages associated with that plugin are removed from the catalog. If a plugin is configured for load on-demand then messages associated with that plugin will be left in the catalog when the plugin is not running. The agent message catalog can be communicated using the agent heartbeat as a list of capabilities.

(74) Also shown is a security manager **320**, which is responsible for configuring security and encryption of messages send to/from the agent **110**, as well as storage and management of encrypted data. The security manager **320** stores related data (e.g., security configuration data and encryption keys) in one or more cryptography files **325**.

(75) Because the data to be sent is stored in the spool file, the agent manager **180** can send data while a plugin is disabled (or enabled).

(76) The agent can be configured dynamically using, for example, Domain Name System (DNS) Service records (SRV records) or a configuration file. In some examples, using DNS SRV records for configuration is preferred when data for a particular DNS domain is sent to a single compliance server. The configuration file setup technique may be preferred when different machines in the same domain will connect to different compliance servers.

(77) In some examples, a provisioning service can be used that informs agents about their initial configuration information, configures the agents with specific combinations of plugins, or provides an upgrade of agent or plugin executable code.

(78) In some examples, a reconciliation service can be used to match previous agent identifiers and operating system information with current identifiers and current operating system information.

This reconciliation service insures continuity in data and logging information stored in the agent data consumers **190**.

(79) Agents can be installed on individual target machines using, e.g., the host operating systems native packaging tools. For example, Windows targets can install agents using the Microsoft MSI installer, while Linux targets can use Red Hat Package Manager (RPM).

(80) Using native system tools allows for easy deployment and upgrade of agents using mechanisms such as Active Directory. The core agent component and associated plugins can each be versioned independently.

(81) On startup the agent collects data including: a list of IP addresses assigned to the host, domain names associated with the list of IP addresses, and performs a lookup for associated DNS SRV records.

(82) D. Agent Plugins

(83) The functionality of agents (e.g., agents **110-112**) can be expanded with the use of agent plugins. For example, FIG. 2 illustrates a number of plugins that are connected to the local agent process through the use of a plugin manager. Plugins can be written in any suitable computer language (e.g., C++) using multi-platform libraries and coding methodologies, which allows for the sharing of the vast majority of agent code across different host platforms. In some examples, other languages besides C++ can be used to implement plugins, provided that support is available for the messaging layer used to connect the agent to its agent bridge. In some examples, Google Protobuf is used as the messaging layer.

(84) As shown in FIG. 3, the agent **110** includes a plugin manager **240** that controls execution of plugins, and routes messages between plugins (e.g., via plugin connectors **245-247**), the message dispatcher **230**, and the spooler **270**. The plugin manager also includes a number of capability maps (e.g., capability map **341**) that list the types of messages and services that are produced and consumed by the connected plugins. The plugin maps can be used to build catalog entries, thereby advertising services available from the agent without exposing additional details regarding the plugins, using, e.g., ConfigRequest messages.

(85) The loading of plugins can be controlled using a rule set, which specifies which plugins to load and connect to the agent, and the order in which to load plugins. In some examples, some plugins can be chained together, thereby providing a degree of modularization.

(86) The plugin manager thread **257** is used to send data to the plugins indicating the current lag of the spooler (or “delta”) to the plugin manager **240** for communication to individual plugins via the plugin connector **245**. Each plugin can have a contract (e.g., an automatically negotiated contract) with the spooler **270** to respond in a certain period of time, and then return to sleep.

(87) In some examples, agents can load plugins on an on-demand basis. In some examples, the agent provides a quarantine functionality to limit resources and/or data that can be accessed by one or more installed plugins. In some examples, plugins can include hooks to allow the plugin to be used across multiple operating system platforms (e.g., by handling both Windows and POSIX hooks).

(88) Plugin Startup

(89) Plugin startup can be initiated based on a number of events, for example, when an agent (e.g., agent **110**) is initialized or after a plugin dies. In some examples, an agent initializes a plugin by issuing a command including arguments sufficient to describe the desired plugin operation (e.g., by passing command line arguments to initialize a plugin process and specifying a path to a configuration file, a path to a log file, read pipe identifier, and/or a write pipe). In some examples, the agent **110** and its associated plugins (e.g., plugins **130-135**) can communicate with each other using other interprocess communications mechanisms provided by the operating system including, but not limited to, shared memory, anonymous pipes, UNIX pipes, streams, or sockets. In some examples, some or all of the plugin startup is delegated to the plugin manager thread **257**, which can also route communications between the plugin manager **240** and the spooler. The configuration

file can include information describing messages that the plugin will receive and send as well as expected behaviors for interacting with the agent that initiated the plugin. Passing arguments at plugin initialization provides the plugin enough information to be fully functional once execution begins.

(90) In some examples, messages sent between agents **110-112**, the agent platform server **150**, and/or the agent data consumers **190** is compressed. Further, plugins can compress data before transferring to its agent, and messages stored in spooler **270** can be compressed. Compressing transmitted messages and/or stored messages can reduce network load and/or reduce the spooler **270** storage capacity requirements. Huffman encoding, Lempel-Ziv encoding, or other suitable compression techniques can be used to compress the messages.

(91) In some examples, executable code for one or more of the plugins can be located in a specified directory on the agent host computer and then discovered and loaded by automatically traversing the specified directory and/or sub-directories for the plugin executables.

(92) In some examples, messages are passed between an agent and its plugins using message envelopes. Messages sent from a plugin can be numbered and addressed by the associated plugin process. The messages can be addressed using a four-part scheme: sender_uuid: the agent identifier from a handshake request message_type: a string value that equates to a specific message type sequence_major_number: a sequence major number provided by the agent sequence_minor_number: a number that can start at an arbitrary value and be increased by the plugin for sending subsequent messages

(93) Together, the sender_uuid, message_type, sequence_major_number, and sequence_minor_number can be used to form envelope addresses. In some examples, a plugin can be added to the agent to handle envelope requirements. By using a major and a minor sequence number, creating file system updates for every message can be avoided, while still providing enough information such that messages can be uniquified in a message spool. This enables agents to find and perform message ResendRequest operations on the spooled messages. Upon invocation of a plugin the agent assigned a new sequence_major_number, thereby making all the messages that the plugin creates, and that the agent writes to the spool, unique even after restart of a plugin or its corresponding agent. An example of major and minor numbers used to address plugin messages are described in Table 1 below.

(94) TABLE-US-00001
TABLE 1 Major Minor Message Number TailConfigResponse 5 1
TailFileData 5 1 . . . 101 Plugin Restart TailFileData 6 1 . . . 99 TailConfigResponse 6 1
TailFileData 6 100 . . .

(95) In some examples of the disclosed technology, a number of different messages types from an agent can be received by its plugins, including handshake requests, status requests, and shutdown requests, as discussed further below.

(96) Handshake Request and Handshake Response

(97) Upon startup, a plugin receives a handshake sent from its corresponding agent. The handshake can include data such as the agent's corresponding identifier (e.g., UUID), a sequence major number, and a data directory identifier. Responsive to receiving the handshake from its agent, the plugin responds with a handshake response message. The handshake response message includes: a plugin identifier (e.g., the name of the plugin), a description of the plugin's capabilities, and a description of messages that will be consumed by the plugin and produced by the plugin. For example, the plugin capabilities list can include a list of message types that can be accepted by the plugin, and the list of plugin messages produced can include configuration and data response message types.

(98) Status Requests

(99) Agents (e.g., agent **110**) can periodically send status messages to one or more of their associated plugins. The plugin in turn responds with a StatusResponse message within a predetermined time period, to notify the agent that the plugin is operating correctly. The status

request can include an indication of the number of seconds between the time (e.g., the wall clock time) of the last message written to the spool and the time of the last message read and sent to the bridge. This provides the plugin with an indication of the state of the agent. For example, if the agent has sent all the messages in its spool, then the message delta is relatively small and sending another small message from the plugin is not expected to burden the agent. Conversely, if the agent has not sent a message to the bridge for some time (e.g., because the agent is disconnected from its bridge, or the agent is behind in sending data to the bridge), then the plugin can choose to buffer more data, thereby creating a larger message before sending to the agent. Sending fewer messages of larger size, (in some examples, up to about 1 Megabyte), incurs less network transmission overhead than sending more messages of a smaller size.

(100) In some examples, plugin StatusResponse messages include the plugin name and a description of its current configuration. This information is collected and stored by the agent manager **180** (e.g., by receiving StatusResponse messages from the plugin manager **240** via the agent's bridge connector **260**). By collecting plugin name and configuration description, a compliance sever (e.g., agent data consumers **190**) can determine whether a particular plugin has an incorrect or outdated configuration, and address the configuration (by, e.g., sending a new plugin configuration to the agent).

(101) Shutdown

(102) In some examples of the disclosed technology, plugins can be shut down as follows. A plugin's host agent sends a Shutdown message to the agent instructing it to shut down. The plugin persists its state (e.g., by storing state information in a computer-readable storage media coupled to the agent) and shuts down. A brief period of time after sending the Shutdown message, the agent closes pipes to/from the agent and sends the plugin's associated process a SIGTERM signal.

(103) ConfigRequest

(104) In some examples of the disclosed technology, the capabilities of plugins can be enhanced using a ConfigRequest to exploit the autonomous nature of agent operations described herein. In a ConfigRequest pattern, a plugin provides a ConfigRequest capability. The plugin responds to ConfigRequests by sending ConfigResponse messages in response, and when configured to do so, additional data with the ConfigResponse.

(105) In some examples of the disclosed technology, a ConfigRequest message for a plugin includes a serial number and a description of a configuration. When a new ConfigRequest is received by a plugin, the plugin replaces its current configuration with the new configuration described in the ConfigRequest. Plugins can store their current configuration and runtime state so that the plugin can resume operation using the current configuration and runtime when the plugin is restarted.

(106) ConfigResponse

(107) The plugin responds with a ConfigResponse message, which tells the Server that a new serial number configuration was received and processed. If the plugin cannot service the entirety of the ConfigRequest, it can include error information in the ConfigResponse. The requested configuration may have been partially applied or rejected completely, as defined by the plugin.

(108) Data Messages

(109) Once a configuration is accepted and applied, the plugin begins sending data messages to the agent data consumers **190** via the agent bridge **160**.

(110) Heartbeat Messages

(111) The agent periodically sends heartbeat messages. The heartbeat messages can contain information including, but not limited to, current wall clock time, current spooler minor/major number, currently-available messages that can be sent to/from plugins for servicing, and/or messages that cannot currently be serviced by any plugin on the agent.

(112) E. Agent Spooler

(113) FIG. 3 further outlines the capabilities and structure of the agent spooler **270**. The agent

spooler **270** includes a number of subcomponents, including a spool file manager **370**, which includes one or more spool files **380**, a message searcher **372** for searching for data in the spool files, and a priority queue **375**.

(114) The spooler stores data as a number of relatively small (e.g., 32 Megabyte) files on disk to form a complete spool (e.g., a 1 Gigabyte spool file comprising 32 each of 32 Megabyte files). Storing spooled data in small files can be desirable, as it limits data loss in the inadvertent event of agent or plugin shutdown or corruption of the smaller 32-Megabyte spool file. Within a spool file, the data can be further segmented in a series of headers and data, where the header indicates a variable length of the data in the segment. As shown in FIG. 3, the spool files **380** are stored in a number of computer-readable storage media and categorized according to the state of the data in the individual files, for example, completed files **381**, pending files **382**, the current read file **383**, and the current write file **384**. Further, markers indicating, e.g., current read and write position of the spool files can be stored in a read position file **385**. As will be readily understood to one of ordinary skill in the art, the spool files **380** are not limited to storage in, for example, a hard drive, but can also be stored in a solid state drive (SSD), non-volatile or volatile memory, a database, or other suitable storage means.

(115) The spool file manager manages the various spool files **381-385** illustrated in FIG. 3. Completed files **381** are spool files that contain messages that have already been sent to the agent bridge **160**. Pending files **382** are spool files that contain unread messages ready to be read and sent to the agent bridge **160**. The current read file **383** is an open spool file that is being read by the bridge writer thread **256** and concurrently forwarded to the bridge connector **260**. After each message is read, the offset of the previously read message and the spool file name is updated in the read position file **385**. This offset and file name lags the last message read by one, thereby enabling recovery if the bridge connector **260** is disconnected from the agent bridge **160** while a message is in flight (being sent). When the connection is reestablished, the message that was in flight is sent again. The agent platform server **150** and the agent data consumers **190** are configured to tolerate duplicate messages. The current write file **384** is the spool file that is open for writing of messages received from the plugins via the message dispatcher **230**. At times the current read file **385** and current write file **383** can be the same file. The spool files **381-385** illustrated in FIG. 3 are one suitable configuration for use with the spooler **270**, but it will be readily understood by one of ordinary skill in the art that other suitable spool file configurations can be employed.

(116) Messages in the spooler **270** can be identified using a unique session identifier (e.g., unique at least for a particular session of operation on a particular agent). In some examples, the unique session identifier includes a major and a minor number. The major number is a unique identifier for the plugin session. The minor number is unique for each message of that type within the corresponding major number run session.

(117) In some examples of the disclosed technology, session identifiers are timestamp agnostic, in other words, the current time on the host computer is not relevant to the ordering of messages according to session identifiers. The session identifiers can be used for relative positioning and sequencing of the messages. In some examples, the spooler **270** can be configured to overwrite a selected portion of the spool (e.g., 20% of the spool). In some examples, there is one spool file per plugin, while in other examples, two or more plugins share the same spool file.

(118) In some examples, message data are encrypted. In some examples, the integrity and authenticity of messages can be verified using HMAC (Hash-based Message Authentication Codes) or other suitable methods to prevent tampering of messages sent by the agent or plugins.

(119) The spooler **270** supports disconnected operations. The supported disconnected operations include spooling when some plugins are disabled and during network disconnections (e.g., intentional or unintentional loss of network connectivity to the agent bridge). By spooling data the agent and its plugins can operate semi-autonomously. Plugins can receive data, configuration information, and perform operations (e.g. vulnerability scans, monitor watch logs, or other

operations) on a regular basis, but the data can be returned regardless of the plugin being enabled/disabled. When the agent **110** reconnects with the bridge **160**, it can send all or a portion of its spooled data one or more of the agent data consumers **190** via the agent bridge **160**.

(120) In some examples, the sizes of the spool file(s) are determined based at least in part on the rate at which messages are being sent. For example, if 0 to 5 messages are being sent, a 100 MB spool file may be sufficient, while sending a larger number of messages may consume a 1 GB or even larger spool file, depending on the amount of data spooled and associated overhead.

(121) Data stored in the spool files can be secured using, e.g., obfuscation by compression, or encryption. In some examples, spool file data is not encrypted, but data in messages sent from the agent to the bridge is encrypted using, e.g., transport layer security (TLS) encryption. The agent should have at least some read access to the spooler data in order to determine, e.g., spool information such as message type, sequence major/minor numbers, and time stamps.

(122) Upon agent startup, the agent spooler **270** creates a new spool file, so that in the event the end of the previous spool file was corrupted, new data will be added to an uncorrupted file, thereby avoiding appending data to a corrupt spool file that may be unreadable.

(123) The priority queue **375** can be used route some messages to the bridge faster, according to priorities assigned to the messages or plugins producing the messages. For example, for a plugin that processes SCAP data, the most recent messages are the most important and thus are desirable to send earlier than lower-priority data. In some examples, all the higher priority messages in the priority queue **375** are sent first, while in other examples, the sending of messages is load balanced so that at least some of the lower-priority messages are sent before the higher priority messages are sent.

(124) In some instances there may be more than one instance of agent spooler **270**. For example, spooler instances can be dedicated to individual plugins, or shared by related plugins.

(125) IV. Example Agent Platform Server

(126) As shown in FIG. 1, the agent platform server provides an agent bridge **160** that can receive messages from agents (e.g., agents **110-112**), a message broker **170** (including one or more message queues **175-178**), and an agent manager. The message broker **170** can be used to route messages between the agent bridge **160** and any of the agent data consumers **190**. Agents (e.g., agents **110-112**) can establish a network connection (e.g., to a number of agent data consumers **190**, including the compliance server **191**, the log server **192**, the policy server **193**, the change management server **194**, etc.) to the agent bridge **160** that is hosted on an agent platform server. The agent connection to the agent bridge **160** can be encrypted using, e.g., Transport Layer Security, Secure Sockets Layer, or other suitable encryption scheme. In TCP/UDP examples, the default port used to connect to the bridge is port number 5670, but other suitable network ports can be employed. To provide additional security, the agent can be configured so that it does not listen for incoming connections from the bridge or compliance server. Instead, the agent initiates communication to these network targets.

(127) Agents can use the underlying operating system routing information to determine how to connect to the agent bridge. Since agents create connections to the bridge, no return routing (from the bridge to an agent) is necessary. Once the connection is established, messages are sent in both directions. For example, an agent bridge can send advisory messages to its associated agents instructing the agents to hold message, or to hold particular types of messages, until a subsequent advisory message is sent indicated that the agent should resume sending of messages (or resume sending particular types of messages). The agents **110-112** and the agent bridge **160** can each establish agent identifiers (e.g., UUIDs) for uniquely identifying the agents and bridges, thereby avoiding reliance on other identifiers, such as IP addresses or MAC addresses that may change or may not provide a unique identifier in virtual environments.

(128) In some operational scenarios one or more agent data consumers **190** may be shutdown for maintenance or is temporarily off-line and undergoing a fault recovery operation. The agent bridge

160 monitors the message queues determining if the agent data consumers **190** are removing messages in a timely fashion. If the message removal is slowing or stopped the agent bridge **160** sends advisory messages to all agents informing them to stop or restart specific message types. This provides a level of fairness for all messages such that a non-operational agent data consumer does not block the traffic of messages to other operational agent data consumers. The message fairness delivery algorithm adds resiliency and robustness in the communication channel from the plugins through the agent and the spooler through the connection to the agent bridge to the message broker and to the agent data consumers. In some examples the agent **110** could have more than one spooler **270** dedicated to a particular plugin or message type facilitating message delivery fairness.

(129) Messages (e.g., messages from any of the agent data consumers **190**) are sent to and from agents via the agent bridge **160**. The agent's bridge subscribes to a topic through which all agent-bound messages travel

(130) Messages can be transported within an envelope. The envelope includes source and/or destination information, and in some examples includes a time stamp, the associated agent identifiers (e.g., UUIDs), message type the sequence major and/or minor numbers, or other information that can be used to route and process the messages contained in the envelope.

(131) In some examples, the agent platform server **150** is hosted on a separate server from any of the agents **110-112**, while in other examples, the agent platform server **150** can reside on the same server as the agents **110-112** and/or the agent data consumers **190**. In some examples, the agent platform server **150** resides in the same local area network (LAN) as the agents **110-112** and/or the agent data consumers, while in other examples, the agent platform server **150** resides at a separate location or in a computing cloud.

(132) A. Agent Bridge

(133) The agent bridge **160** receives and sends messages from and to agents (e.g., agents **110-112**). The agent bridge **160** can notify other components when agent connections **110-112** are created or destroyed, using, e.g., AgentConnect and AgentDisconnect messages. The agent bridge **160** can generate a bridge identifier (e.g., a UUID) for itself that can be used by the agents **110-112** and the agent data consumers **190** to uniquely identify the agent bridge in subsequent messages.

(134) AgentConnect/AgentDisconnect

(135) When an agent connects to the bridge, the bridge creates an AgentConnect message. An agent manager (e.g., executing on one or more of the agent data consumers **190**) subscribes to these messages. Similarly, when an agent disconnects from the bridge, the bridge creates an AgentDisconnect message. The agent manager **180** subscribes to these messages.

(136) B. Agent Manager

(137) In some examples of the disclosed technology, the computing environment **100** includes the agent manager **180**, which can execute on, for example, the agent platform server, or another suitable environment. The agent manager can manage the status of agents and provide information for agent data consumers **190** upon request.

(138) The agent manager **180** sends messages to the agent **110** to configure the agent's associated plugins and processes data received from the agent **110**. Some of the types of agent data that can be tracked by the agent manager include: heartbeats, host names, IPv4/IPv6 addresses, capabilities, capability configurations (capability and serial number), and host platform. In some examples, configuration messages can be initiated by the agent **110** itself, or by one of the agent data consumers **190**.

(139) OnlineAgentsRequest/OnlineAgentsResponse

(140) In some examples of the disclosed technology, the agent manager **180** can also provide OnlineAgentsRequest services. One or more of the agent data consumers **190** sends an OnlineAgentsRequest message to get a list of online Agents. The agent manager **180** responds to this request by sending an OnlineAgentsResponse message. The OnlineAgentsResponse message can be used in conjunction with the ongoing AgentOnline messages to track online agents. The

agent manager **180** sends AgentOnline messages when, for example, an agent connects to the agent bridge or when an agent's capabilities change. In the event that a plugin dies and its associated agent does not restart the plugin, the capabilities for that plugin will no longer be included in catalog published with the associated agents' heartbeat. Upon recognizing the change, the agent manager **180** can publish a new AgentOnline message for the agent. This allows the agent data consumers **190** to discover new plugins or determine that plugins are no longer operational, based on changes in the message catalog reported in the heartbeat messages.

(141) C. Message Broker

(142) The agent platform server **150** can use the message broker **170** for distributing messages between the agent bridge **160**, the agent manager **180**, and the agent data consumers **190**. The message queues **175-178** allow messages to be temporarily stored before sending on to their destination, and can be used to buffer traffic in the event of a failure in the connecting network, one or more of the agents **110-112**, or one or more of the agent data consumers **190**.

(143) V. Example Agent Data Consumers

(144) FIG. **1** illustrates a number of agent data consumers **190**, including a compliance server **191**, a log server **192**, a policy server **193**, a change management server **194**, a file integrity monitoring server **195**, etc. Although a finite number of agent data consumers **190** are shown in FIG. **1**, it will be readily understood to one of ordinary skill in the art that any number of servers can consume agent data, and that some of the agent data consumers **190** can be omitted, depending on the deployment environment.

(145) VI. Example Techniques Performed in Exemplary Agent Systems

(146) A. Disconnected Mode

(147) The agents disclosed herein are designed to be semi-autonomous. For example, agent plugins are designed to accept a complete configuration describing what acts the agent is to perform and when. As a result, plugins can continue operating according to the current configuration without communicating with the agent manager **180** or the agent data consumers **190**.

(148) In other words, an agent and its plugins can continue their normal operations of watching a target machine even while being unable to communicate with the agent bridge **160**, agent manager **180**, and/or agent data consumers **190**. Messages and data generated by the plugin can be spooled by the agent until a connection to the bridge can be re-established and the spooled messages sent.

(149) B. Example Agent Configuration and Registration

(150) FIG. **4** is a diagram **400** that outlines an example of communication between an agent and an agent platform server during agent registration, as can be used in some examples of the disclosed technology. The respective acts and messages performed by an agent platform server and an agent are shown along a timeline. It will be readily understood that this is an example, and that some embodiments of the disclosed technology can add, omit, or rearrange the acts outlined in FIG. **4**. In some examples, the agent is the agent **110** and the agent platform server **150**, as described above, although other agent and agent platform server structures can be employed.

(151) At act **410**, the agent turns Federal Information Processing Standards (FIPS) mode on for use in subsequent communications. At act **411**, the agent checks to determine whether it has already been registered. If the agent has not been registered, then the agent generates keys and a certificate signing request (CSR) at act **412**. Further, at act **412** the agent generates a digital signature and client authorization information. At act **415**, the agent turns FIPS mode off. In other examples, the method proceeds without turning FIPS mode off.

(152) At act **420**, the agent sends a client hello message to the agent platform server. In some examples of the disclosed technology, the agent initiates an anonymous Transport Layer Security (TLS) handshake request with the agent platform server that is further detailed below regarding FIG. **5**. Responsive to receiving the client hello message, the agent platform server sends a server hello message to the agent at act **421**. At act **440**, the agent builds an agent registration request, and sends the request to the agent platform server at act **445**. The agent registration request includes the

keys, CSR, digital signature, and client authorization information that were previously generated. (153) When the agent platform server receives the agent registration request, the server verifies the request. At acts **450** and **451**, the agent platform server verifies the registration key and CSR received with the agent registration request. At act **455**, the agent platform server signs the CSR, and sends the signed CSR to the agent in an agent registration response message at act **460**. (154) At act **470**, the agent converts the signed CSR to a Privacy Enhanced Mail (PEM) format certificate and stores the PEM certificate in a local computer readable storage media. At act **475**, the agent disconnects from the agent platform server.

(155) AgentUUIDChange message

(156) In some examples of the disclosed technology, agents are uniquely identified by an agent UUID, which is 128 bits in length. Some examples of when an agent data consumer **190** can receive an AgentUUIDChange message from an agent include: registration of a new agent, corruption of agent state files, change in one or more MAC addresses on the system, cloning of a virtual machine (VM), and/or replacement of one or more network cards on the agent's host machine.

(157) In the cases of a new agent registration or corruption of an agent state file, an AgentUUIDChange message arrives at the compliance server with a new current UUID and the previous UUID is not set because it did not exist or was unreadable.

(158) In the case of MAC changes, VM changes, and NIC changes, the AgentUUIDChange message can arrive with both the new current and the previous UUID set. In this case, the server should be configured to take appropriate action, e.g., merging agent data if the messages are from the same system, or treating the data as separate in the case of a cloned system. The server needs to determine whether the assets associated with the agent are the same node or different nodes and whether to associate the previous data stream with the new data stream. In some examples, the server is configured to make these determinations automatically, while in other examples, the server receives input (e.g., from a user or a configuration file) to make the determination.

(159) In the event that a virtual machine with the agent is installed is cloned, but not yet started, then the cloned agent will be authenticated (e.g., using password and/or public-key certificates, as described above) and obtain new certificates for connecting to the bridge. The agent will then generate a new UUID as described above.

(160) In the event that a virtual machine with the agent installed is cloned after agent startup, then upon startup, the new agent will connect to the bridge using the existing certificates stored on the cloned system. If all MAC addresses on the agent's host system have changed, the agent will generate a new UUID and send an AgentUUIDChange message with the new and previous UUIDs for the agent. The server will create a new asset record corresponding to the new agent's host system. If not all of the MAC addresses on the agent's host system are changed, then the agent will connect to the bridge using the certificate and existing UUID. This allows for the same agent identifier to be used in, for examples, maintenance scenarios where a network interface card is replaced in a host having more than one network interface card, without generating a new agent identifier.

(161) C. Example Agent Certificate Handshake

(162) FIG. **5** is a diagram **500** that outlines an example of communication between an agent and an agent platform server during agent registration, as can be used in some examples of the disclosed technology. The respective acts and messages performed by an agent platform server and an agent are shown along a timeline. At act **510**, the agent turns FIPS mode on for use in subsequent communication. At act **511**, the agent loads a certificate issued by a certificate authority. In some examples, the certificate authority is CAcert, although other certificate authorities can be used. At act **512**, the agent loads an agent identification certificate. At act **513**, the agent loads an agent privacy key. After loading the certificates at **511-513**, the agent turns FIPS mode off at act **514** and issues a hello message to the agent platform server at act **520**. The hello message includes the

certificates loaded at acts **511** and **512** and is encrypted using the agent privacy key generated at act **514**. In other examples, the method proceeds without turning FIPS mode off.

(163) After receiving the hello message, the agent platform server verifies the agent peer information at act **530**, the certificate at act **531**, and the agent identification certificate at act **532**. Once the information is verified at acts **530-532**, the agent platform server responds to the agent with a hello message at act **540**.

(164) After the agent receives the agent platform server's hello message, it responds by verifying the server peer information at act **550**, the server's certificate authority-issued certificate at act **551**, and the bridge identity certificate at act **552**. Once the information in the server hello message has been successfully verified, the agent responds with an agent heartbeat message **560**. As further detailed above, the agent heartbeat message includes information that can be used to process spooled-off messages from plugins and to identify and request plugin services, including time data, spooler marker numbers, and available messages that are processed by the agent's plugins.

(165) FIG. **6** is a diagram **600** illustrating generation of a certificate **610** by the agent platform server based on both the agent identification certificate **620** loaded at act **512** and the bridge identification certificate **630** sent with the agent platform server hello message at act **540**. The certificate **610** can be used to authenticate communications between agents and agent bridges.

(166) D. Agent/Bridge Server Messaging Flow

(167) FIG. **7** is a diagram **700** that illustrates an example message flow between various agent/bridge components, as can be performed in certain examples of the disclosed technology. As shown in FIG. **7**, a message from the agent bridge **160** (hosted by an agent platform server **150**) is asynchronously read **710** by the bridge connector **260** hosted by an agent (e.g., agent **110**). The bridge connector **260** then dispatches **715** the received message to the message dispatcher **230**. The message dispatcher **230** determines the appropriate plugin for servicing the message, and issues a write message **720** to the appropriate plugin's plugin connector **245**. The plugin connector **245** issues an asynchronous write message **730** to the plugin, and responsive to the message, receives an asynchronous read message **735**. The plugin connector **245** then sends a dispatch message **740** to the message dispatcher **230**, which in turn sends a handle message **750** to the plugin handler **237**. The plugin handler sends a write message **760** to the plugin connector **245** and also sends a write message **770** to the spooler **270**. The data from the plugin is stored in a spool at the spooler **270** until the message can be sent to the bridge by sending a write message **780** to the bridge connector **260**. Once the bridge connector **260** establishes a connection to the bridge, after the plugin services the message, its data is returned to the plugin connector with an asynchronous read. In some instances message writes and reads maybe all be synchronous, asynchronous or a combination of both.

(168) E. First Example Agent Message Sequence: Resend Request

(169) ResendRequest Message

(170) As data is received from one or more agents, the agent data consumers **190** can monitor messages received from the agents (including messages from agent plugins) and request resends of any messages that are missing. In some examples, the bridge (e.g. agent bridge **160**) can perform similar monitoring and request resends. Because the order in which messages are delivered to the agent data consumers (e.g., the compliance server **191**) is not necessarily guaranteed, it is typically advantageous for the compliance server **191** to have handling capabilities for messages received out of order.

(171) One high-level example of how a compliance server handles missing messages is provided below. Assume that a compliance server receives the following stream of data, and that the WelConfigResponse and WelData messages are for the same configuration serial_number.

(172) An example of major and minor numbers used to address plugin messages are described in Table 2 below.

(173) TABLE-US-00002 TABLE 2 Major Minor Message Number Number WelConfigResponse 5

1 WelData 5 1 WelData 5 2 WelData 5 3 . . . 50 WelData 5 55 WelData 5 56 . . . 101 WelData 6 1
(174) As shown in Table 2, a number of messages (5-51 through 5-54) have been dropped.
(175) The first message is a WelConfigResponse message, which is received in response to a previously-sent WelConfigRequest message. The compliance server then starts receiving a number of WelData messages, which have major number 5 and minor numbers, being consecutive integers from 1 to 50. The next WelData message received by the compliance server has the same major number (5) but a minor number of 55, which indicates a gap in WelData for minor numbers 51-54.
(176) Responsive to detecting the missing data, the server sends a ResendRequest message to the agent. An example format for the message is “ResendRequest previous(5, 50),” which indicates that the next five messages starting from minor message number 50 are being requested. The server then waits for data to be sent (e.g., in a ResendResponse message). Assuming the messages were actually dropped, then the missing 4 data messages should be received, followed by a ResendResponse message: status_code FOUND (5, 51) (5, 54)
(177) As will be readily understood to one of ordinary skill in the art, it should be noted that the message names and formatting are for exemplary purposes, and that other suitable formats of messages can be used in certain embodiments.
(178) Rolled Major Numbers Messages
(179) Continuing the exemplary datastream illustrated in Table 2, above, when server receives the WelData messages (5,101) followed by WelData (6,1), it will realize that some data from the agent may be missing. Responsive to determining a gap in the received messages, the server determines the new major sequence number in the received data stream (6) and identifies possible missed messages. The server constructs and sends a ResendRequest for the possibly missing data. The message can be formed as:
(180) TABLE-US-00003 ResendRequest previous(5, 101) next (6,1).
Upon receipt of this ResendRequest message, the agent scan its spool and resends any messages that are identified as missing by the ResendRequest, in a similar fashion to the ResendResponse message discussed above. Conversely, if the agent determines that there are no missing messages to send in response to the ResendRequest, a message is sent to indicate this to the server:
(181) TABLE-US-00004 ResendResponse request = ResendRequest ID status_code NOT_FOUND_MISSING (5,102) (6,0)
(where ResendRequest ID is the identifier for the Resend Request being responded to, and messages 5-102 through 6-0 are the missing messages not found in the spooler).
(182) In some examples, messages that appear to be missing will occur normally, as the major number used by the agent is incremented upon starting or restarting the associated plugin.
(183) Missing or Corrupt Messages
(184) Using the same example of a dropped message discussed above, the agent can also send a number of different status codes, depending on the cause of the missing message(s). Table 3 below lists a number of different status codes that can be used in an exemplary embodiment of the disclosed technology:
(185) TABLE-US-00005 TABLE 3 Code Description FOUND Message was found and resent NOT_FOUND_CORRUPT The internal or data message was corrupt in the spool and messages are/may be missing NOT_FOUND_MISSING The requested message is not in the spool (message may have never existed) NOT_FOUND_TOO_OLD The message is of a range that is before the oldest message in the spool (message may have never existed) NOT_FOUND_UNKNOWN Unknown cause of missing message(s)
(186) The following examples illustrate messages being sent in different scenarios responsive to a ResendRequest.
(187) For example, if the next message received at the server is (major, minor): (5, 100), then the server would expect to receive messages (5,1) through (5,99) previously, otherwise, the messages are determined to be missing. A ResendRequest for messages between (5,1) and (5,100) is sent. If

the agent returns any of the NOT_FOUND_* status codes, then the messages are determined to have been lost.

(188) In another variation of this example, if the next message received is (6,1), then a ResendRequest is sent to determine whether any messages between (5,100) and (6,0) were not received. A ResendRequest for messages between (5, 100) and (6, 0) is sent. If the agent returns the code NOT_FOUND_CORRUPT, then the server determines that data has been lost. Alternatively, if the agent returns the code NOT_FOUND_MISSING, then the server determines that all the data between (5,100) and (6,0) has been received from the agent. Alternatively, if the agent returns the code NOT_FOUND_TOO_OLD, then the result is indeterminate—the server determines that data from the agent may or may not have been lost.

(189) F. Second Example Agent Message Sequence: Resend Request

(190) This section discussed another example of ResendRequest handling, as can be performed in certain embodiments of the disclosed technology. Assume that a system comprising an agent, bridge, and compliance server are in the state depicted in the diagram **800** of FIG. **8**. As shown in FIG. **8**, a number of messages are, or have been, stored in a spool memory **810** (indicated by the files within the dashed lines), which includes a number of individual spool files (e.g., two spooled-off files **820**, four completed files **821**, etc.). Each rectangle depicted within the spool memory **810** represents a different spool file. The spool memory **810** can be implemented within an agent **110**, as described above regarding FIGS. **1-3**, although in other examples, different suitable agent architectures can be employed.

(191) Two spooled-off files **820** have already been removed from the spool and are no longer accessible. The first spooled-off file started at message number (1-10) and ended at (1-25), while the second spooled-off file (1-26) and ended at (1-51).

(192) Four completed spool files **821** have been sent to the agent bridge, but have not yet been removed from the spool. The message data in the completed spool files **821** is available to be re-sent to the agent bridge if requested. As shown in FIG. **8**, the first completed file spans two major numbers, from (1-52 to 1-61 and 2-1 to 2-10). The major number was advanced from 1 to 2 due to, for example, restarting of the agent or restarting of a plugin associated with the spool.

(193) The current read spool file **822** is shown, which includes messages from 3-21 to 3-60. A current read pointer **830** indicates the current read position for spool messages being sent from the agent to the agent bridge. The current read pointer **830** is advanced as additional messages from the spool are sent to the agent bridge.

(194) Two pending files **823** are shown, which are queued to be read after the current read file **822** has been completely read and send to the agent bridge.

(195) The current write file **824** is shown, which starts at position 3-105 and currently ends at 3-127. The current write pointer **835** stores the current position for writing within the current write file **824**.

(196) As shown in FIG. **8**, each of the spool files **820-824** includes one of the time stamps **840** shown. The time stamps can be used to determine the current lag time between messages. Lag time information can be used to, for example, adjust the rate at which messages are sent to the agent bridge or the rate at which plugins produce data.

(197) In some examples of the disclosed technology, each of the spool files **820-824** includes a header indicating the size of the spool file, an envelope, and message data. Messages used to transmit data from the agent to its agent bridge can include an envelope with data that can be used to identify the agent and/or the plugin producing the message data, for example: message type, agent UUID, destination UUID, major/minor number, and timestamp can be included in the message envelope.

(198) As an example, assume that the system of FIG. **1** is being used to spool plugin data to the agent data consumers **190**. One of the agent data consumers **190** determines that some agent data may be missing, and the agent data consumer sends the agent **110** a ResendRequest message **850**.

The ResendRequest message **850** includes an indication of the previous message (1-41) and next message (3-32) before and after the messages that may be missing, thereby indicating a range of messages to be searched. The ResendRequest message **850** also includes the respective time stamps (20140202081200 and 2014020607007) for the possible missing messages.

(199) The agent receives the ResendRequest message **850** and searches the spool memory **810** (e.g., using a message searcher **372**). The agent message search determines that: messages 1-42 through 1-51 are too old and have been deleted from the spool (“TOO_OLD”); messages 1-52 through 1-61, 2-01 through 2-54, and 3-01 through 3-31 were found and are still stored in the spool (“FOUND”), and messages 1-62 through 2-00 and 2-55 and 3-00 are missing (“MISSING”) (because the messages were never generated). The agent sends a ResendResponse **860** indicating the results of the agent message search to one or more of the agent data consumers **190**. In some examples, the agent also automatically sends the FOUND messages to the agent data consumer that sent the ResendRequest, while in other examples, the agent waits for an additional request to send particular messages from the spool.

(200) The agent data consumers **190** that receive the ResendResponse **860** can react accordingly. For some applications, lost or missing data may merely be noted, while in other applications, the data can be used to, for example, initiate additional vulnerability scans to replace the missing data.

(201) G. Server/Agent/Plugin Contract

(202) In some examples of the disclosed technology, there is the shared contract for server components, agents, and their plugins. The contract can be expressed in an Interface Definition Language (IDL) and implemented using an Application Programming Interface (API) to coordinate between server components, agents, and their plugins. The contract establishes a syntax for creating properly formed messages, including defining required and optional fields in the messaging protocol.

(203) Under the example contract, the agent is responsible for full verification of a plugin package before it is considered and is available to be launched. The plugin package includes: executable files, configuration information, and/or command line arguments including a manifest. The plugin package can be digitally signed using suitable cryptographic techniques. Agents (e.g., agents **110-112**) can verify the plugin's digital signature before the plugin is considered for operation. In some examples, the plugin package is re-verified every time the plugin is started, while in other examples, the plugin package is re-verified periodically, at system startup, or other suitable intervals. Plugins can connect to the agent using only the pipes given to it from the command line arguments. Plugins use the configuration and log directories passed to it from the command line arguments. The first message an agent sends to a plugin after connect is a HandshakeRequest message. The agent includes full path information for plugin's directory and the plugin's executable directory (plugin package directory). The executable directory is used for locating shared libraries the plugin may use or delegate to another application, for example: a Real-Time Manager application (RTMGR). A RTMGR application can communicate to operating system kernel modules to obtain user-specific security information that is associated with a particular event or change or operation that a plugin is monitoring.

(204) The first message from a plugin sent to the agent is a HandshakeResponse message after it has received a HandshakeRequest. The plugin uses the data directory given to it in the HandshakeRequest it receives from the agent. The Agent periodically sends StatusRequest messages to all plugins. Plugins respond to StatusRequest messages from the agent with StatusResponse messages.

(205) Plugins are expected to respond immediately to a Shutdown request message. The plugin periodically persists its last working position if the plugin cannot immediately persist when it receives a Shutdown message from the agent.

(206) Upon restart, plugins resume work from the state using its last ConfigRequest. The plugin manages any incomplete units of work.

(207) If a plugin has no work to do and it is designed for on-demand loading, it informs the agent by sending a PluginExiting message. The agent subsequently moves the plugin's capabilities to the on-demand maps and sends a Shutdown message to the Plugin and removes the plugin from the plugin list. The agent is responsible to re-launch the plugin if a message is received that is contained in the message catalog. Servers, agent and plugins operate to keep the amount of message data in the spooler at or below a selected limit (e.g., to avoid using more than 80% of the spooler capacity). In some examples, the selected limit for the current capacity can be adjusted. Agent and plugin messages can be formed and enveloped as Protobuf messages. Thus, the agent does not need to have detailed information about the plugin's structure or operation.

(208) H. Agent Authentication:

(209) In some examples of the disclosed technology, the bridge is configured to use a registration key for authentication. When in this registration mode, the agent must supply the correct registration key to the bridge upon first connecting to the bridge in order to authenticate to the bridge. The initial key can be sent using anonymous SSL. After successfully authenticating the key, the agent subsequently receives an encryption certificate (e.g., a public-key certificate such as an X.509 certificate) that can be used for encrypting messages sent using subsequent connections to the bridge. If the registration key is changed (e.g., by system administrator re-configuring the bridge) the agent can continue to use the certificate that was received early using the older key. However, any new agent instances will need to use the new registration key to acquire their corresponding certificates.

(210) VII. Exemplary Method of Spooling of Host Data

(211) FIG. 9 is a flow chart 900 illustrating an example method of spooling host data, as can be performed in some examples of the disclosed technology. For example, an agent hosting an agent spooler 270 in the computing environment 100 discussed above can be employed to perform the illustrated method.

(212) At process block 910, an agent operating on a computing host having one or more network connections collects host data from, e.g., one or more plugins on the computing host. The agent is configured to collect the host data whether or not the agent can currently send data via any of the network connections. The types of data collected using the plugins can include FIM data, COCR data, WEL data, Windows Registry data, or other suitable data. After collecting the data, the method proceeds to process block 920.

(213) At process block 920, the agent receiving a message from an agent bridge indicating a message type to send to at least one of the agent data consumers. Based on receiving the message, the agent initiates sending of data for the indicated message type. The agent determines whether it can send data via the host computer's network connection(s), and if so, the method proceeds to process block 930. If not, the method proceeds to process block 940.

(214) At process block 930, the agent has determined that it can send data via the network, and the agent proceeds to send at least a portion of the spooled host data to at least one of agent data consumers. For example, data can be sent to any of the agent data consumers 190 illustrated in FIG. 1 via an agent bridge 160. In some examples, the collected data is temporarily stored in the spooler before being sent.

(215) At process block 940, the agent stores at least a portion of the collected host data in a spooler for later transmission. In some examples, the spooler 270 discussed above can be employed. In some examples, some of the spooled data can be overwritten or removed based on the priority of the data, or according to other suitable criteria. In some examples, messages of a first message type are stored in a first spool of the spooler and messages of a second message type are stored in a second spool of the spooler. In some examples, the rate at which spooled data is sent to the agent data consumers is increased or decreased based at least in part on lag of the spooler.

(216) In some examples, the rate at which data is spooled and/or the rate at which the spooled data is sent to the one or more agent data consumers is based at least in part on currently-available

resources of the host computer. In some examples, one or more plugins executable on the host computer collect the host data, the agent sends data to the plugins indicating lag of the spooler, and the plugins adjust the rate of collecting host data based at least in part on the indicated spooler lag. In some examples, the host data is collected by one or more plugins executable on the host computer when the agent cannot send data via the network connection, and the spooled host data is sent to the at least one of the agent data consumers via a bridge executing on an agent platform server.

(217) After the collected host data is sent (according to process block **930**) or stored (according to process block **940**), the method can proceed to collect more host data. As the computer network connection becomes available or unavailable, the agent can elect to store data in the spooler or send data accordingly. The rate at which the spooled data is sent can be increased or decreased based on, for example, the lag of the spooler or host computer resources. In some examples, the data that is collected is based on a request for a type of data. The agent searches for plugins that can provide the requested type of data, and invokes a corresponding plugin if found. In some examples, agent can send a description of the types of data and/or messages that can be produced by the agent.

(218) In some examples of the method, a second message is received from an agent bridge indicating a second message type to not send to at least one of the agent data consumers. Based on the receiving the second message, the agent stores for messages of the second message type in a spooler until receiving a third messaging indicating the spooled messages of the second type are to be sent.

(219) VIII. Exemplary Method of Sending Messages with Agent-Generated Sequence Numbers

(220) FIG. **10** is a flow chart **1000** illustrating an example method of resending messages based on agent-generated sequence numbers, as can be performed in some examples of the disclosed technology. For example, a computer hosting an agent **110** operating in the computing environment **100** discussed above can be employed to perform the illustrated method.

(221) At process block **1010**, an agent executing on a host computer, sending one or more data messages to a server with the computer network. The data messages including sequence numbers generated by the agent. For example, each of the sequence numbers can include a major number, which is incremented upon starting or restarting the agent, and a minor number, which is incremented with each message sent. Further examples of suitable sequence numbers as can be used in some examples of the disclosed technology are discussed above and illustrated in FIG. **8**.

(222) At process block **1020**, the agent receives a resend message from the server indicating that one or more of the data messages are to be resent. The messages to be resent are indicated using at least some of the sequence numbers, for example, a range of sequence numbers can be used. The ResendRequest message **850** illustrated in FIG. **8** is an example of suitable resend message. After receiving the resend message, the method proceeds to process block **1030**.

(223) At process block **1030**, the agent searches for the indicated messages based on the one or more of the generated sequence numbers. In some examples, the messages may have already been removed from the agent's host computer. After searching for the messages, the method proceeds to process block **1040**.

(224) At process block **1040**, the agent can resend one or more of the requested messages and resend any messages still stored on the agent's host computer. In some examples, an additional reply message, such as the ResendRespond message **860** illustrated in **860**, can be sent with the resent data messages to describe which of the requested messages have been found, are too old, or are missing.

(225) IX. Exemplary Method of Identifying Agent Messages

(226) FIG. **11** is a flow chart **1100** illustrating an example method of identifying agent messages sent on a computer network, as can be performed in some examples of the disclosed technology. For example, a computer hosting an agent **110** operating in the computing environment **100** discussed above can be employed to perform the illustrated method.

(227) At process block **1110**, the agent **110** self-generates a unique agent identifier for itself. The agent identifier is independent of any network addresses associated with the agent's host computer. The agent identifier can be, for example, a UUID. After generated the agent identifier, the method proceeds to process block **1120**.

(228) At process block **1120**, the agent sends a first message to at least one agent data consumer (e.g., one or more of the agent data consumers **190**). The agent data consumer can stored the agent identifier for use in determining the origin of subsequent messages from the agent. After sending the first message, the method proceeds to process block **1130**.

(229) At process block **1130**, the agent is moved to a different physical and/or virtual host computer. The agent then sends a second message to the agent data consumer including the same agent identifier. By using the same identifier, the agent data consumer can track messages from the same agent even though the agent is operating from a different network address.

(230) In some examples, of the exemplary method, the agent can be replicated (e.g., on one or more additional physical and/or virtual hosts), and each of the replicated agents in turn self-generates a unique agent identifier for itself, thereby allowing agent data consumers to distinguish amongst the replicated agents and the original agent. Further, the same agent identifier can be used in scenarios where, e.g., the agent's network address changes when a network interface card is changed, but without changing the agent identifier.

(231) X. Exemplary Message Delivery Fairness Methods and Apparatus

(232) FIGS. **12-16B** illustrate methods and apparatus for improved message delivery fairness between agents and servers using a computer network, as can be practiced in some examples of the disclosed technology. Disclosed message delivery fairness methods enable continued delivery of messages from agents, even though one or more consuming servers are offline for maintenance or are undergoing fault recovery operations. For example, the exemplary computing environment **100** described in further detail above can be used to perform the exemplary methods outlined in each of FIGS. **12-16** using the described agents (e.g., agents **110-112**) and their associated plugins **130-135**, agent platform server **150** and associated agent bridge **160**, message queues and/or topics, and agent manager, and agent data consumers **190** (e.g., agent data consumers **191-198**). As will be readily understood by one of ordinary skill in the art, the disclosed methods can also be practiced in other suitable computing environments that have been adapted for use with one or more of the disclosed methods.

(233) FIG. **12** is a diagram **1200** that illustrates an exemplary set of message flows between agents (e.g., agent **1210**) and agent data consumers **1220** using a computer network. As shown, the message flow is parallel between a number of agent plugins **1230** and their associated agents; serial through the agent and spool subsystem to an agent bridge **1240**; and parallel again once the messages are placed in the message broker server topics. In scenarios where one of the consuming servers is offline for maintenance or is undergoing a fault recovery operation, unconsumed messages will queue in the inbound topic and will eventually fill the message broker's inflight message holding capacity. When a particular broker is full, some or all of the message flow through the broker will stop. In turn, all operational message consumers become starved, as the bridge is not able to place any new messages on the broker's message topics. Subsequently, this stops the message flow from all connected agents. Since the message flow through the connected agents is serial, the spools will eventually fill, causing the agents to shut down their plugins, and hence message collection and flow to the servers will be stopped.

(234) Two factors that can affect message flow between agents and agent data consumers includes (1) the total inflight message capacity of the broker-once the message capacity is reached, message flow through the broker stops; and (2) the serial message flow region of agent input to the spooler through the bridge connection(s) blockage of any message can cause a blockage of all messages behind it, and can eventually force the agent to shut down plugins once the spool is full.

(235) One way to compensate for finite message capacity of a broker is to allow the unconsumed

messages of a topic to be deleted if there is no active consumer. This can be undesirable in some application, because it can cause thrashing throughout the system as a flood of resend requests and response operations are created, and this can lead to unrecoverable loss of data.

(236) In some examples of the disclosed technology, serial message flow between agents and brokers can be modified to have parallel channels of multiple spoolers (one per message family or per plugin) and add additional, parallel transport layer security (TLS) connections between an agent and its bridge. However, adding additional connections to the bridge may be undesirable, because of the substantial increase in committed server resources for the bridge, which reduces its connected agent capacity. Using dedicated spoolers per plugin can shorten the serial path between plugins and bridges.

(237) In some examples of the disclosed technology, a mechanism is included for the agent bridge **1240** to communicate a message instructing agents to stop sending message types that are blocked in the message broker, and to send only unblocked messages to the bridge. When the blocked message types are flowing again, the Bridge sends a message instructing the Agents to start sending messages again. Having separate spools for plugins can thereby simplify implementation of this mechanism in the agent.

(238) FIG. **13** illustrates is a diagram **1300** illustrating an exemplary message flow, including unidirectional merging, as can be used in some examples of the disclosed technology.

(239) As shown in FIG. **13**, an agent bridge **1310** monitors the message topics **1320** to the consuming data servers **1330**. The agent bridge **1310** sends advisory messages to some or all of its connected agents if the number of messages in a message topic reaches a configurable limit, dubbed a “Stop Sending Level,” which indicates that the consuming server has reached its limit (e.g., because the consuming server is offline, having problems consuming messages at the inbound message rate, or has another issue that causes the limit to be reached). A message topic is a form of a queue, where messages can be sent to more than one consumer (e.g., messages from agents can be queued and then sent to each of a plurality of agent data consumers). If the limit is met or exceeded an advisory message is sent instructing agents to stop sending messages of the associated type(s). Once the issue with the consuming server is resolved (e.g., the server is back online) and the topic spool is reduced down to a lower configurable limit (dubbed “Restart Sending Level,” an advisory message is sent to the connected agents instructing them to start sending the appropriate message types. Also shown is an auxiliary bridge **1315** that can be used to send similar messages as the agent bridge **1310** when, for example, the agent bridge **1310** is disabled or unreachable, to offload message load from the agent bridge **1310**, or to send messages to particular topics and/or data consumers based on an associated message type (e.g., as allocated by an affinity service).

(240) Scenarios in which the agent bridge can be configured to send advisory messages include the following. For example, after an agent has connected to the agent bridge and validated its credentials, the agent bridge sends the agent the most recent valid advisory message. In some examples, the advisory message is only sent to newly-connected agents.

(241) Advisory messages can be sent under varying system conditions. For example, the bridge can send advisory messages immediately after an agent has connected and credentials have been validated. In response, the bridge sends to the newly-connected agent the most recent valid advisory message. In this case, the message is only sent to the newly connected agent. The bridge can also send advisory messages to multiple, or all, connected agents when one or more topics message count exceeds the stop sending level or drops below the restart sending level.

(242) The agents **1340** depicted in FIG. **13** have multiple spoolers **1350** (e.g., one spooler per plugin **1360**, **1361**). The agents **1340** listen for advisory messages from the bridge **1310**. The agents can send or skip over message types based on the advisory messages.

(243) FIGS. **14**, **15A** and **15B**, and **16A** and **16B** are sequence diagrams that illustrate message flows: under normal operation, during the sending of a stop sending advisory message, and during the sending of a restart sending advisory message, respectively.

(244) FIG. **14** is a sequence diagram **1400** that illustrates an example message flow during normal operation of an agent with two active plugins, dubbed WEL (a Windows event log plugin) and FIM (a file integrity monitoring plugin), as can be practiced in some examples of the disclosed technology. The agents are connected to the bridge and successfully sending messages that are consumed by agent data consumers. As will be readily understood by one of ordinary skill in the art, the example messages flows are not limited to WEL and FIM plugins, but can be readily adapted for use with other suitable agent plugins.

(245) FIGS. **15A** and **15B** depict a sequence diagram **1500** that illustrates an example message flow during a stop sending operation, as can be practiced in some examples of the disclosed technology. An agent with two active plugins, WEL and FIM, is connected to the bridge. In this example, the FIM Server (an agent data consumer) is taken down for maintenance. The bridge determines that the FIM message topic has reached the Stop Sending Level, and in response sends Advisory messages instructing the connected agents to stop sending FIM messages.

(246) FIGS. **16A** and **16B** depict a sequence diagram **1600** that illustrates an example message flow during a restart sending operation, as can be practiced in some examples of the disclosed technology. An agent with two active plugins, WEL and FIM, is connected to the bridge. In this example, the FIM server has been restarted, and starts to receive messages from the FIM message topic. As a result, the bridge determines that the count of messages in the FIM topic has dropped below the restart sending level, and in response sends advisory messages to agents instructing them to restart sending FIM messages.

(247) In some examples of the disclosed technology, the bridge is configured so that it will not stop completing agent connection requests in the event message topics in message broker are blocked. In some examples of the disclosed technology, it may be undesirable to bridge message queue level oscillate between the stop sending level and the restart sending level. In some examples, so a closed loop lag controller is employed to use a recent sample history and a computed threshold that is exceeded before the advisory messages are sent.

(248) In examples including multiple spools, a number of tools can be used to manage the multiple spools. For example, specific plugin spools can be allocated a predefined percentage of the overall spool size. In some examples, learning techniques are used to observe the proportions in which message flows are occurring, and to adjust spool allocation based on actual proportions. In some examples, a combination of predefined percentage and learning techniques can be applied to manage multiple spool allocation.

(249) In some examples, the bridge has additional support for a mix of agents that do or do not support message delivery fairness techniques.

(250) In some examples, agents wait for a slight delay between completing a connection to the bridge and/or waiting for an initial connect advisory message before, before the agent begins reading data from its associated spools.

(251) In some examples of the disclosed technology, a single common spool is used for messages from all plugins of an agent. If advisory messages are received from the bridge, the agent can write messages that the bridge has advised against sending to temporary spool files. These temporary spool files contain messages of a single type, in order in which they were read from the main spool. When the bridge sends advisory messages to start sending, the agent temporarily stops reading data from the main spool and reads the messages in the temporary spool files, sending the temporary spool file messages to the Bridge. Once all the temporary spool files are read, the temporary spool files are erased and the Agent resumes reading from the main spool. By prioritizing the reading of the temporary spool files, the agent allows for keeping the message sequencing temporarily correct for each message type.

(252) XI. Example Affinity Service

(253) FIG. **17** is a block diagram illustrating an example computing environment **1700** in which an affinity service can be deployed according to the disclosed technology. For example, the disclosed

affinity services can be used with the agents **110-112**, agent platform server **150**, and agent data consumers **190** discussed above with respect to FIG. **1**.

(254) A number of agents **1710**, **1711**, and **1712** are illustrated in FIG. **17**. Similar to the agents **110-112** described regarding FIG. **1**, each of the agents **1710-1712** can include a local agent process that can manage and communicate with a number of plugins, including a file integrity monitoring (FIM) plugin and a log plugin (e.g., a Windows event log (WEL) plugin) that are configured to extend the functionality of the respective agent.

(255) Each of the agents **1710-1712** communicates with other components in the system depicted in the computing environment **1700** via an agent platform server **1750**. As shown, the agent platform server **1750** includes an agent bridge **1760** for sending messages to and from agents (e.g., agents **1710-1712**). The agent bridge **1760** can send messages over a computer network to agents executing on other computers, using inter-process and/or inter-thread communication to agents executing on the same computer as the communication bridge, or by using other suitable communication means. The illustrated agent platform server **1750** also includes a message broker **1770** with one or more message queues for temporarily storing messages received from and sent to, for example, the agent bridge **1760**, an agent manager **1780**, an affinity service **1785**, and agent data consumers **1790**. The agent platform server **1750** coordinates operation of the agents by sending and receiving messages using the message broker **1770**. In some examples, the agent platform server **1750** is configured to send an advisory message to adjust the rate at which messages are sent by one or more agents. As shown in FIG. **17**, the affinity service **1785** resides as a component of the agent platform server **1750** (e.g., as a standalone process executing on the agent platform server **1750**), while in other examples, the affinity service is hosted in an alternate location (e.g., as a thread or other component of the agent manager **1780**).

(256) In some examples of the disclosed technology, for example, in large networks with multiple agent platform servers **1750** and multiple agent data consumers **1790**, the affinity service **1785** would be external to the agent platform server and centralized to improve communications with all instances of the agent platform server and agent data consumers.

(257) As shown in FIG. **17**, the agent data consumers **1790** include multiple log servers (**1790-1** and **1790-2**) and multiple FIM servers (**1795-1**, **1795-2**, and **1795-3**). In some examples the multiple log servers and/or the multiple FIM servers are hosted on separate virtual machines on the same physical hardware (e.g., a computing server). In some examples, the multiple log servers and/or the multiple FIM servers are hosted on separate physical machines in the same computer network environment. In some examples, multiple log servers and/or the multiple FIM servers are hosted on separate physical machines in different computing environments.

(258) The affinity service **1785** provides mappings to the message broker **1770** and/or agent bridge **1760** in order to direct message flow from the agents (e.g., agents **1710-1712**) to one of the multiple log servers and/or multiple FIM servers. The affinity service **1785** can utilize UUIDs in order to identify the agents **1710-1712** and agent data consumers **1790**.

(259) In some examples, the affinity service **1785** maintains a table representing the associations between agents (e.g. agents **1710-1712**) and one or more of the agent data consumers **1790**). The agents can be assigned using a number of methodologies, including but not limited to assignments based on: round robin, load and/or capacity of one or more of the agent data consumers **1790**, geographic location of the agents and/or the agent data consumers, network topology (e.g., by physical subnets or virtual local area network (VLAN)), function roles (e.g., a respective consumer and/or agent is deployed for product development, testing, staging, or production), version of an agent, and/or version of an agent data consumer.

(260) In some examples, the affinity service **1785** directs routing of messages from agents by intercepting an agent online message emitted by the agent manager **1780**. The agent online message is enhanced by providing the product server UUID assigned to the agent by the affinity service **1785**.

(261) In some examples, the affinity service **1785** maintains an affinity map that defines relationships between agents and agent data consumers. In some examples, the affinity service is configured to map each of the agents to a respective one of the data consumers. In some examples, the affinity service mapping is based at least in part on one or more of the following: a geographic location of one or more of the agents and/or the agent data consumers; topology of a network carrying communication between the agent data consumers, agent platform servers, and/or agent computing hosts; a functional role of one of the agents and/or one of the agent data consumers; a version of an agent; and/or a version of an agent data consumer.

(262) FIG. **18** is a sequence diagram **1800** illustrating an example of messages sent between an agent, bridge, manager, affinity service component, and data consumer, in an exemplary method of assigning an agent to a data consumer using the affinity service component, as can be performed in some examples of the disclosed technology. For example, hardware in the example computing environment discussed above regarding FIG. **17** can be used to implement the illustrated components.

(263) In the illustrated example, affinity assignment based on a round robin allocation is used, although other examples of allocation may be used in other examples. Upon initial agent discovery by an agent manager (in this example, an agent manager executing a thread for the affinity service), the affinity service hosted by agent manager will determine which of the data consumers is to be assigned to an agent. The agent manager issues an AgentAffinity message, which includes information about the assignment, including UUID and an AgentOnline message. In some examples, the affinity service sends the AgentAffinity message to a specific data consumer. In other examples, the AgentAffinity message is broadcast to all or a portion of the data consumers, and the non-addressed consumers drop or ignore the message (e.g., if the UUID does not match the data consumer's assigned UUID).

(264) In examples where the AgentAffinity message is for a specific instance of an application (e.g., a CCC tool), the CCC tool can create a new node entry for the newly-associated agent. Thus, consumers can ignore subsequent AgentOnline messages that are not stored for a previously-created node, and update the node if the AgentOnline message for an agent associated with a particular instance of a consumer.

(265) An example of an affinity map, before any agents have been assigned by an instance of the affinity service **1785**, is illustrated as a text listing **1900** in FIG. **19**. An example of an affinity map, after a number of agents have been assigned to agent data consumers using the affinity service **1785**, is illustrated a text listing **2000** in FIGS. **20A-20B**. As illustrated in FIGS. **19-20B**, affinity maps define relationships used to associate agents and agent data consumers and can be used, for example, to route messages from the agents to the agent data consumers. Examples of hierarchical information that can appear in an affinity map include a map definition, including a unique name for the affinity map, a map type (e.g., a map type could be for one type of agent, but other agents and/or components can have different assigned types, a region (e.g., a geographical region or role), a version identifier, a valid time (e.g., a field describing valid time period for the affinity map). Additional information can include server definitions, including server type, identification of a method used to assign affinity, a capabilities list (e.g., describing capabilities specified for an agent to be associated, and/or a description of capabilities for the associated agent data consumer), a unique name, list of agents, and/or a list of associated agent UUIDs.

(266) XII. Example Method of Controlling Message Flow in a Computer Network

(267) FIG. **21** illustrates an example method of controlling message flow in a computer network that couples a plurality of agents, a plurality of agent data consumers, and an agent message bridge configured to send messages between the agents and the agent data consumers, as can be performed in certain examples of the disclosed technology. For example, the computing environment **1700** discussed above regarding FIG. **17** can be used to implement the illustrated method.

(268) At process block **2100**, an agent bridge receives a set of messages from one or more agents,

at least some of the messages including a message type for the type of data carried by the message.

After the agent bridge receives the messages, the method proceeds to process block **2110**.

(269) At process block **2110**, a set of messages are queued in a spooler. For example, messages containing agent data to be sent to the agent bridge can be spooled including an indication of a respective message type for each of the messages. In some examples, the spooler includes distinct spools for queuing the set of messages, and each message of the set of messages is placed in a selected one of the distinct spools based on the type of the message.

(270) At process block **2120**, one or more of the agents receives an indication that sending of one or more, but not all of the messages in the spooler should be delayed for at least one indicated message type. For example, messages of a type to be sent to an agent data consumer that is experiencing network difficulties, high loading, reconfiguration, or other such conditions can be indicated in order to delay sending of agent messages of the indicated type to the agent data consumer. In some examples, the indication is received based on a queue depth reported by a message broker. In some examples, the indication is received from the agent bridge responsive to determining that one or more of the agent message consumers are not receiving messages.

(271) At process block **2130**, at least one message of another type than the type indicated at process block **2120** is sent to an agent data consumer via the message bridge. Thus, messages that are not indicated to be delayed using the spooler can be sent, for example, to a different agent data consumer than the consumer associated with the indicated message type. In some examples, the selected agent data consumers to which to send messages are determined using an affinity service.

(272) In some examples of the illustrated method, the agent bridge later sends an indication that the sending of the delayed messages queued in the spooler can be resumed, and the respective agents proceed to send messages based on spooler data to one or more agent data consumers for the indicated message type.

(273) In some examples, the illustrated method further includes mapping each of a plurality of the agents to a respective one of the agent data consumers, the mapping being based on at least one or more of the following: a geographic location of one or more of the agents and/or the agent data consumers; topology of a network carrying communication between the agent data consumers, agent platform servers, and/or agent computing hosts; a functional role of one of the agents and/or one of the agent data consumers; a version of an agent; and/or a version of an agent data consumer.

(274) In some examples, the agent bridge monitors a plurality of message topics or queues, each of the message topics or queues having a distinct type, each of the message topics or queues being configured to temporarily store messages received from the agents. When the number of messages queued in a first message topic or queue of the plurality of message topics or queues exceeds a predefined stop sending level, the bridge sends an advisory message from the agent bridge to the agents indicating that messages of the corresponding type should not be sent. After sending the stop sending advisory message, when the number of messages queued in the first message topic or queue reaches a predefined restart sending level, the agent bridge sends an advisory message from the message bridge to the agents indicating that sending of messages of the corresponding type can be resumed.

(275) XIII. Example Computing Environment

(276) FIG. **22** illustrates a generalized example of a suitable computing environment **2200** in which described embodiments, techniques, and technologies, including reporting agents and monitor servers, can be implemented. For example, the computing environment **2200** can implement any of the agents, agent platform servers, and agent data consumers, as described herein.

(277) The computing environment **2200** is not intended to suggest any limitation as to scope of use or functionality of the technology, as the technology may be implemented in diverse general-purpose or special-purpose computing environments. For example, the disclosed technology may be implemented with other computer system configurations, including hand held devices, multiprocessor systems, microprocessor-based or programmable consumer electronics, network

PCs, minicomputers, mainframe computers, and the like. The disclosed technology may also be practiced in distributed computing environments where tasks are performed by remote processing devices that are linked through a communications network. In a distributed computing environment, program modules may be located in both local and remote memory storage devices. (278) With reference to FIG. 22, the computing environment **2200** includes at least one central processing unit **2210** and memory **2220**. In FIG. 22, this most basic configuration **2230** is included within a dashed line. The central processing unit **2210** executes computer-executable instructions and may be a real or a virtual processor. In a multi-processing system, multiple processing units execute computer-executable instructions to increase processing power and as such, multiple processors can be running simultaneously. The memory **2220** may be volatile memory (e.g., registers, cache, RAM), non-volatile memory (e.g., ROM, EEPROM, flash memory, etc.), or some combination of the two. The memory **2220** stores software **2280**, images, and video that can, for example, implement the technologies described herein. A computing environment may have additional features. For example, the computing environment **2200** includes storage **2240**, one or more input devices **2250**, one or more output devices **2260**, and one or more communication connections **2270**. An interconnection mechanism (not shown) such as a bus, a controller, or a network, interconnects the components of the computing environment **2200**. Typically, operating system software (not shown) provides an operating environment for other software executing in the computing environment **2200**, and coordinates activities of the components of the computing environment **2200**.

(279) The storage **2240** may be removable or non-removable, and includes magnetic disks, magnetic tapes or cassettes, CD-ROMs, CD-RWs, DVDs, or any other medium which can be used to store information and that can be accessed within the computing environment **2200**. The storage **2240** stores instructions for the software **2280**, plugin data, and messages, which can be used to implement technologies described herein.

(280) The input device(s) **2250** may be a touch input device, such as a keyboard, keypad, mouse, touch screen display, pen, or trackball, a voice input device, a scanning device, or another device, that provides input to the computing environment **2200**. For audio, the input device(s) **2250** may be a sound card or similar device that accepts audio input in analog or digital form, or a CD-ROM reader that provides audio samples to the computing environment **2200**. The output device(s) **2260** may be a display, printer, speaker, CD-writer, or another device that provides output from the computing environment **2200**.

(281) The communication connection(s) **2270** enable communication over a communication medium (e.g., a connecting network) to another computing entity. The communication medium conveys information such as computer-executable instructions, compressed graphics information, video, or other data in a modulated data signal. The communication connection(s) **2270** are not limited to wired connections (e.g., megabit or gigabit Ethernet, Infiniband, Fibre Channel over electrical or fiber optic connections) but also include wireless technologies (e.g., RF connections via Bluetooth, WiFi (IEEE 802.11a/b/n), WiMax, cellular, satellite, laser, infrared) and other suitable communication connections for providing a network connection for the disclosed agents, bridges, and agent data consumers. In a virtual host environment, the communication(s) connections can be a virtualized network connection provided by the virtual host.

(282) Some embodiments of the disclosed methods can be performed using computer-executable instructions implementing all or a portion of the disclosed technology in a computing cloud **2290**. For example, agents can be executing vulnerability scanning functions in the computing environment while agent platform (e.g., bridge) and agent data consumer service can be performed on servers located in the computing cloud **2290**.

(283) Computer-readable media are any available media that can be accessed within a computing environment **2200**. By way of example, and not limitation, with the computing environment **2200**, computer-readable media include memory **2220** and/or storage **2240**. As should be readily

understood, the term computer-readable storage media includes the media for data storage such as memory **2220** and storage **2240**, and not transmission media such as modulated data signals. (284) In view of the many possible embodiments to which the principles of the disclosed subject matter may be applied, it should be recognized that the illustrated embodiments are only preferred examples and should not be taken as limiting the scope of the scope of the claims to those preferred examples. Rather, the scope of the claimed subject matter is defined by the following claims. We therefore claim as our invention all that comes within the scope of these claims.

Claims

1. At least one computer-readable storage medium storing computer-executable instructions, which when executed by a processor in communication with a server, cause the server to perform a method, the computer-executable instructions comprising: instructions that cause the server to send a message to at least one agent executing on a compute host indicating a message type of a plurality of message types via a communication connection; instructions that cause the server to send advisory messages to the at least one agent instructing to hold all messages having one or more message types of the plurality of message types while still forwarding all messages having a message type other than the one or more message types, wherein the one or more message types are not all of the plurality of message types; instructions that cause the server to send a subsequent advisory message indicating that the agent should resume sending of messages having at least one of the one or more message types; and instructions that cause the server to receive host data from the at least one agent, wherein at least some of the received host data has been stored at the compute host when the at least one agent cannot send data to the server via the communication connection.
2. The computer-readable storage medium of claim 1, wherein the computer-executable instructions further comprise: instructions that cause the server to send the host data to at least one agent data consumer.
3. The computer-readable storage medium of claim 1, wherein the computer-executable instructions further comprise: instructions that cause the server to send the received host data to at least one agent data consumer selected based on an agent identifier indicated by a message received from the host computer.
4. The computer-readable storage medium of claim 1, wherein the computer-executable instructions further comprise: instructions that cause the server to monitor one or more message topics; and instructions that, when a number of message topics reaches a configurable limit, cause the server to send an advisory message to one or more agents indicating to stop sending or restart sending of messages according to at least one of the message topics.
5. The computer-readable storage medium of claim 1, wherein the computer-executable instructions further comprise: instructions that cause the server to send a message to the at least one agent to adjust the rate at which messages are sent to the server.
6. A server, comprising: at least one processor; a communication connection operable to communicate with the at least one agent executing on a host computer; and a computer-readable storage medium storing computer-executable instructions, which when executed by the at least one processor, cause the at least one processor to perform a method, the computer-executable instructions comprising: instructions that cause the at least one processor to send a message to at least one agent executing on a compute host indicating a message type of a plurality of message types via a communication connection; instructions that cause the at least one processor to send advisory messages to the at least one agent instructing to hold all messages having one or more message types of the plurality of message types while still forwarding all messages having a message type other than the one or more message types, wherein the one or more message types are not all of the plurality of message types; instructions that cause the at least one processor to send a

subsequent advisory message indicating that the agent should resume sending of messages having at least one of the one or more message types; and instructions that cause the at least one processor to receive host data from the at least one agent, wherein at least some of the received host data has been stored at the compute host when the at least one agent cannot send data to the server via the communication connection.

7. The server of claim 6, further comprising: a message broker means for a temporarily storing the received host data.

8. The server of claim 6, wherein the computer-executable instructions further comprise: instructions that cause the at least one processor to send the host data to at least one agent data consumer.

9. The server of claim 6, wherein the computer-executable instructions further comprise: instructions that cause the at least one processor to send the received host data to at least one agent data consumer selected based on an agent identifier indicated by a message received from the host computer.

10. The server of claim 6, wherein the computer-executable instructions further comprise: instructions that cause the at least one processor to monitor one or more message topics; and instructions that, when a number of message topics reaches a configurable limit, cause the at least one processor send an advisory message to one or more agents indicating to stop sending or restart sending of messages according to at least one of the message topics.

11. The server of claim 6, wherein the computer-executable instructions further comprise: instructions that cause the at least one processor to send a message to the at least one agent to adjust the rate at which messages are sent to the server.

12. A method comprising: communicating, by at least one processor in communication with a server, with at least one agent executing on a host computer; and sending, by the at least one processor, a message to at least one agent executing on a compute host indicating a message type of a plurality of message types via a communication connection; sending, by the at least one processor, advisory messages to the at least one agent instructing to hold all messages having one or more message types of the plurality of message types while still forwarding all messages having a message type other than the one or more message types, wherein the one or more message types are not all of the plurality of message types; sending, by the at least one processor, a subsequent advisory message indicating that the agent should resume sending of messages having at least one of the one or more message types; and receiving, by the at least one processor, host data from the at least one agent, wherein at least some of the received host data has been stored at the compute host when the at least one agent cannot send data to the server via the communication connection.

13. The method of claim 12, further comprising: temporarily storing, by the at least one processor, the received host data.

14. The method of claim 12, further comprising: sending, by the at least one processor, the host data to at least one agent data consumer.

15. The method of claim 12, further comprising: sending, by the at least one processor, the received host data to at least one agent data consumer selected based on an agent identifier indicated by a message received from the host computer.

16. The method of claim 12, further comprising: monitoring, by the at least one processor, one or more message topics; and when a number of message topics reaches a configurable limit, sending, by the at least one processor, an advisory message to one or more agents indicating to stop sending or restart sending of messages according to at least one of the message topics.

17. The method of claim 12, further comprising: sending, by the at least one processor, a message to the at least one agent to adjust the rate at which messages are sent to the server.
