

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12395452
Kind Code	B1
Date of Patent	August 19, 2025
Inventor(s)	Biran; Yahav et al.

Scaling-aware traffic request prioritization for scalable connection-based resource

Abstract

A scaling-aware prioritization manager for a scalable connection-based resource receives request traffic and monitors the traffic. The manager may also monitor a connection capacity to the connection-based resource or request processing capacity of the scalable connection-based resource. The prioritization manager enables priority-based routing of new request traffic, based on a determination that the scalable connection-based resource is at or within a threshold limit of a maximum scale capacity for the scalable connection-based resource and disables the priority-based routing of new request traffic, based on a determination that the scalable connection-based resource is no longer at or within a threshold limit of the maximum scale capacity. The prioritization manager may provide an interface for customers to configure the prioritization. The prioritization manager may include an artificial intelligence component for detecting anomalies and generating or modifying priority policies.

Inventors:	Biran; Yahav (Redmond, WA), Katz; Jonathan S. (New York, NY)
Applicant:	Amazon Technologies, Inc. (Seattle, WA)
Family ID:	1000007169711
Assignee:	Amazon Technologies, Inc. (Seattle, WA)
Appl. No.:	18/192586
Filed:	March 29, 2023

Publication Classification

Int. Cl.: H04L47/70 (20220101); H04L47/74 (20220101); H04L47/83 (20220101)

U.S. Cl.:

CPC H04L47/828 (20130101); H04L47/748 (20130101); H04L47/83 (20220501);

Field of Classification Search

CPC: H04L (47/828); H04L (47/748); H04L (47/823)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7673057	12/2009	Yip	370/232	H04L 47/822
10021008	12/2017	Pai	N/A	H04L 47/748
2011/0069685	12/2010	Tofighbakhsh	370/235	H04L 43/0858
2019/0173804	12/2018	Nicas	N/A	H04L 47/822
2019/0340033	12/2018	Ganteaume	N/A	H04L 43/16
2020/0026579	12/2019	Bahramshahry	N/A	G06F 9/5077
2021/0211391	12/2020	Paraschiv	N/A	G06F 9/3877
2023/0318988	12/2022	Wang	709/226	H04L 47/83

OTHER PUBLICATIONS

AWS, “Amazo Aurora User Guide for Aurora”, Retrieved from <https://docs.aws.amazon.com/pdfs/AmazonRDS/latest/AuroraUserGuide/aurora-ug.pdf#rds-proxy>, Updated on Mar. 22, 2023, pp. 1-2156. cited by applicant

AWS, “Connection management and pooling”, Retrieved from <https://docs.aws.amazon.com/whitepapers/latest/amazon-aurora-mysql-db-admin-handbook/connection-management-and-pooling.html>, 2023, pp. 1-2. cited by applicant

Primary Examiner: Chacko; Joe

Attorney, Agent or Firm: Kowert, Hood, Munyon, Rankin & Goetzel, P.C.

Background/Summary

BACKGROUND

(1) Availability of connection-based resources (e.g., resources of a computing system such as a compute resource service, database resource service, or other type of storage resource service, or combinations thereof) can be limited. Clients (e.g., network-based applications such as websites, various processes, or the like) that rely on such connection-based resources can be interrupted when the connection-based resource is unable to process all of the requests from the client in a timely fashion. In a database example, a single database writer may be used to host different client application use cases. If the resources of that database server are exhausted, the client application may be interrupted. For example, the connection-based resource (the database server) may drop requests from the client application indiscriminately, or at least not in a manner coinciding with the business logic associated with the client application. Such indiscriminate interruptions hinder client application processes, such as transactional flows that unexpectedly break in the middle due to the indiscriminate drops by the connection-based resource.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is a block diagram that illustrates a logical architecture of a system for scaling-aware traffic request prioritization for a scalable connection-based resource, according to some embodiments.
- (2) FIG. 2 is a block diagram that illustrates a logical architecture of a system, with a connection pooler, for scaling-aware traffic request prioritization for a scalable connection-based resource, according to some embodiments.
- (3) FIG. 3 is a process diagram that illustrates a process for scaling-aware traffic request prioritization for a scalable connection-based resource, according to some embodiments.
- (4) FIG. 4 is a process diagram that illustrates a request processing process for a connection pooler instance within the context of scaling-aware traffic request prioritization for a scalable connection-based resource, according to some embodiments.
- (5) FIG. 5 is a process diagram that illustrates a process for scaling connection poolers and a connection-based resource, using scaling-aware traffic request prioritization, according to some embodiments.
- (6) FIG. 6 is a process diagram that illustrates a scaling process for a connection pool within the context of scaling-aware traffic request prioritization for a scalable connection-based resource, according to some embodiments.
- (7) FIG. 7 is a block diagram that illustrates an example service provider network embodiment for scaling-aware traffic request prioritization for a scalable connection-based resource, according to some embodiments.
- (8) FIG. 8 illustrate an example of a computer system, one or more of which may implement various components associated with scaling-aware traffic request prioritization for a scalable connection-based resource, described and illustrated throughout the disclosure, according to embodiments.
- (9) While the solution is described herein by way of example for several embodiments and illustrative drawings, those skilled in the art will recognize that the solution is not limited to the embodiments or drawings described. It should be understood, that the drawings and detailed description thereto are not intended to limit the solution to the particular form disclosed, but on the contrary, the intention is to cover all modifications, equivalents and alternatives falling within the spirit and scope of the present solution. Headings used herein are for organizational purposes only and are not meant to be used to limit the scope of the description.

DETAILED DESCRIPTION OF EMBODIMENTS

- (10) As discussed in more detail below, systems and methods for embodiments of scaling-aware traffic request prioritization for a scalable connection-based resource are disclosed. In an example embodiment, request-generating Clients (e.g., processes or applications, etc.) generate and send requests to a Connection-Based Resource (e.g., a backend service or process or similar). A Request Priority Manager, between the Clients and the Connection-Based Resource, receives the request traffic from the Clients and routes the requests to particular instances of the Connection-Based Resource (over established connections between the Request Priority Manager and the Connection-Based Resource, in embodiments).
- (11) In embodiments, a Request Priority Manager may start and stop (enable/disable or apply/halt, etc.) prioritization of processing of traffic requests, based on scaling information (e.g., a scaling state or capacity) of the Scalable Connection-Based Resource. For a Scalable Connection-Based Resource that is near-to or approaching a maximum scaling limit, based on either a configuration setting (e.g., a configurable maximum scaling limit or threshold) or based on physical limits of the resources processing the request traffic) for example, the Request Priority Manager may start prioritizing portions of the traffic (e.g., by type or otherwise) such that lower-priority traffic is throttled (e.g., delayed or dropped, etc.) so that higher-priority or critical portions of the traffic

requests are processed. Such a Request Priority Manager may facilitate improved processing over other systems because the Request Priority Manager continues to successfully process high-priority or critical traffic requests despite the Scalable Connection-Based Resource approaching, or reaching, a point where processing of the entirety of the traffic requests may be delayed or fail in an indiscriminate fashion.

(12) In embodiments, a Request Priority Manager may scale various resources such as connections carrying request traffic to and/or from the Request Priority Manager and/or may instruct the target Scalable Connection-Based Resource to scale. The Request Priority Manager may start prioritizing portions of the traffic over the connections, such as but not necessarily limited to, when either the connections are nearing maximum scale capacity or when resources of the Connection-Based Resource are nearing a maximum scale capacity.

(13) The Request Priority Manager may be scaling-aware in various different ways, and use information about scale (e.g., scaling information or state, or responses to requests indicating a maximum scale has been reached) to enable/disable prioritized routing of request traffic. For example, if the connections over which the traffic is transmitted (e.g., some or all of the various connections from the Clients to the target Connection-Based Resource) have reached a maximum scale, or if the target Connection-Based Resource has reached a maximum scale, and traffic continues to grow, overwhelming the maximum scale of components of the system, there is risk of dropped or delayed requests/responses. Without the priority-based routing described herein, such drops/delays may be performed indiscriminately, or at least in a manner that does not necessarily align with business objectives associated with the requesting Client. In embodiments, the Request Priority Manager may, based on one or more scaling characteristics of the system, determine to enable priority-based traffic request routing. For example, a Client may specify a priority of various different types of requests (e.g., based on Client-based priorities, reflecting business logic, in embodiments) based on characteristics of the requests (such as based on the target object identified in database request, or based on a source ID of a request, as a few non-limiting examples). Such priority may be specified by the client via a priority policy, for example. The Request Priority Manager **110** may route the requests in accordance with the specified priority when some scaling metric has reached some threshold value, for example, and then return to normal, performing (non-priority) routing to the Connection-Based Resource when some scaling metric falls back from some threshold value (when some part of the system is no longer at, near, nor approaching maximum scale).

(14) In one non-exhaustive example used throughout the Detailed Description, a database service is used as one example of a Connection-Based Resource, but embodiments of other services are also contemplated, without limitation. The database service (e.g., a transactional database or other, using a connection pooler or not) hosts a database server on behalf of clients or customers across one or more servers (e.g., single or multi-master and read replica databases). Requesting Clients, such as an application (e.g., such as but not limited to a microservices application) may use a database server that requires many database connections to insert, update and read database objects. The database can host multi-tenant applications to save on hosting dedicated databases, in a non-limiting example.

(15) Continuing with the example, if a connection pooler is used, the connection pooler may instruct or request for new database instances to be created, based on client needs for example (e.g., when the database writer instance starts to run out of resources). For example, a single database writer may be used to host different application use cases, and if the resources of that database server are exhausted, the application service can be interrupted (e.g., transactional flows can break in the middle, an unsatisfactory outcome). In embodiments, the connection pooler scales and contracts autonomously based on client demand. Existing database connections may be reused to reduce database anthropogenic effects caused by volatile client demand.

(16) At least some embodiments described herein utilize a Request Priority Manager including or

separate from a feature such as a connection pooler endpoint (load balancer-based or otherwise) in front of a Connection-Based Resource. In an embodiment wherein the Connection-Based Resource is a database service, the Request Priority Manager may be in front of the database writers and readers to manage communications between Requesting Clients (e.g., users, applications, etc.) and the database servers.

(17) In embodiments, the Request Priority Manager (e.g., including a connection pooler, or otherwise) ensures that enough database resources are available and may request for resources of the database to scale as needed (e.g., bring on more databases readers or writers, for example). In embodiments, the Request Priority Manager prioritizes critical requests of the Connection-Based Resource (e.g., in the case of a database service, such as but not limited to database objects (tables, views and users)) when available resources of the database cross thresholds to maximize application availability and throughput. In an example, the Request Priority Manager may receive (or intercept, in some architectures) requests directed to the Connection-Based Resource (e.g., database requests) and route the requests based on characteristics of the requests it prioritizes (e.g., based on the database objects it prioritizes, according to a priority policy (sometimes referred to as a prioritization policy)).

(18) In a non-exhaustive example, in the case where a Request Priority Manager works in concert with a connection pooler to route incoming requests targeting a Connection-Based Resource such as a database service, the pooler can maintain a list of potential Connection-Based Resource instances (e.g., database instances). If there are none available, the connection pooler can request a new instance. The connection pooler receives requests (e.g., like \$dbuser or \$object). The Request Priority Manager can prioritize specific \$user or \$objects requests (read or write) based on pre-defined policy. For example, the Request Priority Manager may prioritize a critical path of an application: facilitating placement of new orders and getting the orders delivered, in the case of an order fulfillment micro service and/or corresponding objects such as orders, deliveries, etc. The Request Priority Manager may determine an anomalous pattern (e.g., performance-based or security-based anomalies, etc.) in the request traffic (e.g., based on use of an artificial intelligence service or component to identify or determine the pattern) and modify a respective priority policy to change a prioritization for new request traffic corresponding to the anomalous pattern. The Request Priority Manager may issue alerts when the database usage is anomalous to mitigate database hacks caused by Requesting Clients such as an application outside of the database.

(19) In another non-exhaustive example, providing a description of a data flow in the context of a Request Priority Manager for a scalable connection-based resource (e.g., a database) an application opens a connection to the database and issues a request (e.g., a \$dbuser set/get a new \$object (e.g., orders)). In the example, the request is routed via a network load balancer (transport layer (TCP)) to a traffic Request Priority Manager with a connection pooler instance. The connection pooler receives the request (get the \$dbuser and \$object) and determines whether it has enough resources to serve the request. If the connection pooler has enough resources, it routes the request to one of the available database instances (e.g., insert/update/delete requests go to the database writer, while selects go to one of the database read replicas).

(20) If the connection pooler(s) is/are at some threshold level (e.g., such as but not limited to 75% of maximum connection pooler capacity) the connection pooler instance spawns more instances like itself (e.g., such as but not limited to, increase 10% of the current connection pooler fleet). Such spawning may eventually trigger database writer scaling, in embodiments. If the maximum scale of connection poolers has been reached, the Request Priority Manager can enable priority-based routing, dropping unnecessary request traffic, and passing only critical request traffic (\$dbuser and \$object) to the database, in one example. Over time, Client demand may fall (request traffic decreases), traffic prioritization may end (all traffic may pass through) and as the connection pooler instance becomes idle, the connection pooler instances are reduced in number.

(21) In embodiments, the connection poolers (or the Request Priority Managers) publish metrics to

a centralized datastore (e.g., a logging and monitoring management tool or service) and issue alerts when anomalous usage patterns are detected. In some embodiments, the Request Priority Manager determines and sends security alerts based upon anomalies in network traffic patterns caused by abnormal resource usage.

(22) In some embodiments, priority-based routing may be implemented as a filter that is always enforced, in accordance with the prioritization features specified in one or more priority policies. Enforcement may mean that routing of the requests is performed in accordance with a priority policy that specifies normal routing if a maximum scaling threshold has not been reached and priority-based routing when a maximum scaling threshold is exceeded, for example.

(23) At least one technical problem addressed by some of the disclosed embodiments of systems and techniques described herein is that clients (e.g., network-based applications such as websites, various processes, or the like) that rely on connection-based resources can be interrupted when the connection-based resource is unable to process all of the requests from the client in a timely fashion. In a database example, a single database writer may be used to host different client application use cases. If the resources of that database server are exhausted, the client application may be interrupted. For example, the connection-based resource (the database server) may drop requests from the client application indiscriminately, or at least not in a manner coinciding with the business logic associated with the client application. Such indiscriminate interruptions hinder client application processes, such as transactional flows that unexpectedly break in the middle due to the indiscriminate drops by the connection-based resource, for example.

(24) Embodiments described herein provide a technical solution to some such problems and constitute an improvement to computer-related technology. For example, embodiments of scaling-aware traffic request prioritization for a scalable connection-based resource utilize a Request Priority Manager in front of database writers and readers to manage communications between client-based users (e.g., applications or other types of request-making client) and database servers such that, even as traffic transmission and/or processing reaches and/or grows beyond a maximum scale of the respective resources, prioritized traffic is more likely to be transmitted and processed. For example, for the case where the Request Priority Manager has a Connection Pooler, the Connection Pooler can scale and contract autonomously based on client demand; and existing database connections can be reused to reduce database anthropogenic effects caused by volatile client demand. In some such architectures the technical benefits may include one or more of: (1) ensuring that enough database resources (or other type of back-end resources) are available and triggered to scale as needed (2) prioritizing critical requests of database objects (tables, views and users) (or other types of backend resources) when available database resources cross maximum scaling thresholds to maximize application availability and throughput and (3) triggering performance-based or security-based alerts based upon anomalies in network traffic patterns caused by abnormal resource usage.

(25) Attention is now brought to the Figures. The following description begins by describing logical components of FIG. 1 that may perform functionality in FIGS. 3 and 4 and then goes on to describe logical components of FIG. 2 that may perform functionality of FIGS. 3-6, although, components from different block diagrams may perform functionality from different process diagrams, in embodiments. Generally, FIGS. 1, 2, 7 and 8 illustrate logical components of a system or device(s), one or more of which may perform some or all of the features illustrated in FIGS. 3-6, in various combinations. For example, the Request Priority Manager 110 may perform the scaling-aware traffic request prioritization techniques illustrated in FIGS. 3-4, for a Scalable Connection-Based Resource 130. It is contemplated that other components than those illustrated may perform some of the disclosed functionality, in some embodiments, and/or that various of the disclosed functionality may be performed in a different order than that described in the example embodiments provided herein. For example, a system or device may perform scaling-aware traffic request prioritization techniques illustrated in FIGS. 3-4 with or without a connection pooler, or at

least some of the functionality illustrated in FIGS. 3-6 and described herein may be performed by a portion of a Request Priority Manager (e.g., Request Priority Manager **110**) co-located with a Connection Pooler **220** and a portion of the Request Priority Manager, or another separate Request Priority Manager (e.g., Request Priority Manager **210**) co-located with the Scalable Connection-Based Resources. In some embodiments, Request Priority Manager **210** may include some or all of the features illustrated for Request Priority Manager **110** (e.g., one set of priorities may be applied at Request Priority Manager **110**, while another set of priorities may be applied at Request Priority Manager **210**, for example). A service-provider, network-based embodiment is then described, in accordance with FIG. 6, components of which may perform some or all of the functionality illustrated in process diagrams 3-6, for example.

(26) It is contemplated that in at least some embodiments, functionality associated with the Request Priority Manager **110** may be performed at a frontend component (e.g., at a frontend database engine or similar). In some embodiments, the Scalable Connection-Based Resource **130** may include resources of a backend component (e.g., database read/write instances) without limitation.

(27) FIG. 1 is a block diagram that illustrates a logical architecture of a system for scaling-aware traffic request prioritization for a scalable connection-based resource, according to some embodiments. In FIG. 1, Clients **102A**, **102B**, . . . , **102Z** send request traffic **120** to a connection-based resource (e.g., illustrated as Scalable Connection-Based Resource **130**). A Request Priority Manager **110** is illustrated and performs routing of the traffic **120** in-between the Clients **102** and the Connection-Based Resource **130**. In the illustrated embodiment, Request Traffic **120** may include individual requests, sent over connectionless links, or over relatively short-lived links, while the Connections **140** between the Request Priority Manager **110** and the Scalable Connection-Based Resource **130** may be relatively longer-lasting connections, over which groups of the requests are transmitted (e.g., maintained by a Connection Pooler, in embodiments). The Request Priority Manager **110** routes the request traffic **120**. In embodiments, the Request Priority Manager **110** sometimes routes the traffic normally, without prioritization (e.g., when the connection-based resource and connections **140** are sufficient to process all of the traffic) and at other times, routes only prioritized traffic (e.g., when the connection-based resource and connections **140** are insufficient (or are nearing a threshold where they begin to become insufficient) to process all of the traffic).

(28) In embodiments, the priority may be specified by the client. For example, a customer of a service provider may specify (e.g., via a priority policy) prioritization of the traffic in a manner that coincides with customer business logic. Using such policies to limit the routing to requests that have a higher priority over other requests may increase the likelihood that higher priority requests are not indiscriminately dropped or delayed. Providing client-specification of the priority for requests facilitates client-based control for determining which requests are dropped or delayed.

(29) FIG. 1 illustrates Request Priority Manager **110** that receives request traffic **120** for a Scalable Connection-Based Resource **130** from clients **102A-Z**. A Connection-Based Resource Scale Capacity Monitor **116** monitors a scaling-based capacity (e.g., a connection capacity or request processing capacity of the Scalable Connection-Based Resource **130**). In response to a determination that the Scalable Connection-Based Resource **130** is at, near-to or within a threshold limit of a maximum capacity for the Scalable Connection-Based Resource, the Request Priority Manager **110** determines to begin applying a priority policy (e.g., obtained from, or according to a policy stored in, Priority Policy Data Store **118**) to newly received request traffic for the Scalable Connection-Based Resource **130** to prioritize some requests over others. Note that the Request Priority Manager **110** can stop applying the priority policy to newly received request traffic for the scalable connection-based resource in response to a determination that the Scalable Connection-Based Resource is not at or within another threshold limit of the maximum capacity for the scalable connection-based resource (e.g., as illustrated in FIG. 3, described below).

(30) In embodiments, the priority policy (e.g., specified by a client and stored in Priority Policy

Data Store **118**) specifies characteristics of the requests (e.g., types of or IDs for data objects stored at the Scalable Connection-Based Resource **130**; requestor identifiers (IDs) etc.) and/or a corresponding priority for the characteristic. The Request Priority Manager **110** may apply the priority policy to prioritize requests having the characteristics (e.g., requests that access the data objects specified in the priority policy or that originate from a source corresponding to, or are sent on behalf of, the requestor ID).

(31) In a further example of prioritization based on a target object, clients may generate, configure, or otherwise determine priority policies for certain tables. For example, tables containing time-sensitive or critical data (e.g., order tracking for an ecommerce site) may be given a higher priority than other tables such as, but not limited to logging for auxiliary, non-critical functions such as garbage collection processes or the like. By using priority policies that give the requests targeting the tables containing time-sensitive or critical data a higher priority over auxiliary, non-critical functions the system increases the likelihood that the time-sensitive or critical data are processed in a timely fashion, thereby helping to maximize client application availability and throughput.

(32) Different priority policies may be specified for different Clients **102**, for different types of traffic (e.g., database traffic vs. micro-service-based traffic, etc.).

(33) Scalable Connection-Based Resource **130** is illustrated with Resource Instance Scaling Manager **132**, which may manage and/or control scaling of the Resource **130**. In some embodiments, Resource Instance Scaling Manager **132** may receive scaling instructions from Request Priority Manager **110**.

(34) FIG. 2 is a block diagram that illustrates a logical architecture of a system, with a Priority-Aware Scalable Connection **220**, for scaling-aware traffic request prioritization for a Scalable Connection-Based Resource **130**, according to some embodiments. FIG. 2 illustrates that the Request Priority Manager **110** includes a Connection Pooler **220**. The Connection Pooler **220** may, for example, establish front-end connections with request-making Clients **102A-Z** for receiving the request traffic **120** for the Scalable Connection-Based Resource **130** and maintain a pool of backend connections **140** to the Scalable Connection-Based Resource **130** for sending the request traffic **120** to the Scalable Connection-Based Resource **130**. For example, a Connection Pooler Instance **222A** may receive connection requests from Clients **102A** and **102B**, and route traffic resulting from those connection requests over an already-existing-at-the-time-it-was-received connection **140** to the target Connection-Based Resource **130**. The Connection Pooler Instance **222A** may continue to maintain the Connection **140** to the Connection-Based Resource **130**, even after connections to the requesting Clients **102A**, **102B** have ended (e.g., to service further client connection requests). FIG. 2 illustrates Connection Pool Scale Manager **228**, that may manage scaling of the Connection Pooler Instance **222A**, **222B**, . . . , **222W**, in response to increases or decreases in the request traffic **120**, for example.

(35) In at least some embodiments, Request Traffic **120** from Clients **102A-Z** may pass through and be load-balanced by a Load-Balancer **230** prior to reaching the Priority-Aware Scalable Connection Pooler **220**.

(36) Priority-Aware Scalable Connection Pooler **220** is illustrated with Connection Pooler Instances **222A-W**, each including a Cache **224** and Connection Router **112** and Traffic Priority Component **114**. In embodiments, Caches **224A**, **224B**, . . . , **224W** are used by respective Connection Pooler Instances **222A-W** to cache local copies of corresponding priority policies. Generally, connection pooler instances are assigned traffic from particular ones of the Clients **102A-Z** (e.g., by Connection Pool Scale Manager **228**) and aggregate the requests. Respective Connection Routers **112A**, **112B**, . . . , **112W** route the requests to respective Resource Instances **134A-X** over respective connections **140**. When a (locally-cached) priority policy is enabled (e.g., by Traffic Priority Component **114A**) respective traffic directed via Connection Pooler Instance **222A** is prioritized in accordance with the corresponding priority policy, by a respective Traffic Priority Component **114A**, for example.

(37) FIG. 2 illustrates a Priority Policy Data Store **118** for storing client-specified policies that describe how the requests are to be prioritized, when enabled. In the example, to apply a priority policy the Request Priority Manager **110** determines, according to the corresponding priority policy, which requests of the request traffic **120**, received via front-end connections from the Clients **102A-Z**, are sent over the backend connections **140** to the Scalable Connection-Based Resource **130**.

(38) In embodiments, responsive to an increase in the request traffic **120** for the Scalable Connection-Based Resource **130**, the Connection Pooler **220** may request additional backend connections **140** to the Scalable Connection-Based Resource **130** to increase a number of connections in the pool of backend connections **140**. In some embodiments, the determination that the Scalable Connection-Based Resource **130** is at or within the threshold limit of the maximum capacity is based on the ability of the Scalable Connection-Based Resource **130** to accept additional backend connections **140**.

(39) FIG. 2 illustrates that the Scalable Connection-Based Resource **130** includes a group of Resource Instances **134A, 134B, . . . , 134X** that establish the backend connections **140** with the Connection Pooler **220** and that process requests from the request traffic received via respective connections of the backend connections **140**. In embodiments, the Scalable Connection-Based Resource **130** is scalable to increase or decrease a number of Resources Instances **134A-X** in the group of resources instances. In embodiments, the Request Priority Manager **110** requests that the Scalable Connection-Based Resource **130** increase the number of Resource Instances **134** in the group of resources instances (e.g., in response to an indication that the Scalable Connection-Based Resource **130** is unable to provide more additional backend connections via the existing resource instances).

(40) In the illustrated embodiment, FIG. 2 illustrates Request Priority Manager **210** located at the Scalable Connection-Based Resource **130**. It is contemplated that in various embodiments, independent Request Priority Managers may be implemented in middleware between the Clients **102** and Connection-Based Resource **130** (e.g., as part of load balancer, as part of a router, etc.) or as part of the Connection-Based Resource **130**, or that a distributed Request Priority Manager may be implemented in multiple parts, as illustrated in FIG. 2, enabling various different types of prioritization (e.g., prioritization based on source of request, prioritization based on target storage object indicated in the request), at different network points, for example. In FIG. 2, the Request Priority Manager **210** receives request traffic for the Scalable Connection-Based Resource **130** routed over the backend connections **140** by the Connection Pooler **220** and monitors the connection capacity or request processing capacity of the Scalable Connection-Based Resource **130**.

(41) The Request Priority Manager **110** is illustrated with a Resource Scale Capacity Monitor **116** for monitoring the Connection-Based Resource **130**. In some embodiments the Resource Scale Capacity Monitor **116** may determine the scale and/or track the scaling state of the Connection-Based Resource **130** based on responses to requests, by the Request Priority Manager **110**, to add connections or additional resources, or may interact directly with the Connection-Based Resource **130** (e.g., via API or otherwise) to determine the scale and/or track the scaling state of the Connection-Based Resource **130**.

(42) In embodiments, in response to a determination that the Scalable Connection-Based Resource **130** is at or within a threshold limit of a maximum capacity for the Connection-Based Resource, the Request Priority Manager **210** begins applying another priority policy to newly received request traffic for the Scalable Connection-Based Resource **130**, in order to prioritize some requests over others. In at least some embodiments, the other priority policy is different from the priority policy applied or enabled by Request Priority Manager **110**. The policy applied by Request Priority Manager **210** may be stored local to the Connection-Based Resource **130** or at Priority Policy Data Store **118**, in various embodiments.

(43) FIG. 2 illustrates Policy Interface **218**, a public-facing, customer-facing, or administrative-level interface (e.g., application program interface (API), command-line interface (CLI), graphical user interface (GUI) or other interface) via which characteristics of requests that are used to group the requests into request types (e.g., based on one or more characteristics of the requests, such as based on the source of the request, based on a target of the request, based on a security level of the request, based on whether the request falls into an anomalous category, etc.) and via which priorities of the requests (priority of the request types) may be specified. The Policy Interface **218** may be used to generate new policies and/or update policies, in embodiments.

(44) FIG. 2 illustrates a Request Traffic Artificial Intelligence Component **221** that is illustrated with Anomaly Detector **223** and Priority Generator **225**. In embodiments, Anomaly Detector may detect anomalies in the request traffic (e.g., via a combination of machine learning and/or artificial intelligence techniques applied to the traffic through the Connection Pooler Instances **222A-W** or applied to metrics from a metrics and/or logging and monitoring service (described with regard to FIG. 6, below) about the traffic or about performance of system components, etc.). For example, data used by the artificial intelligence techniques may be obtained via local monitoring and logging agents at the Connection Pooler Instances **222A-W**. The Traffic Anomaly Detector may obtain metrics about the traffic from the local agents to create a larger, global picture of the traffic, in order to identify anomalies in the overall traffic, in embodiments.

(45) In embodiments, the machine learning or artificial intelligence techniques may make use of historical request traffic data, known patterns, and the like to identify anomalies and generate and transmit alerts indicating the anomalies.

(46) Priority Generator **225** may monitor whether traffic is being routed in a prioritized manner. In some embodiments, Priority Generator **225** may increase a priority of non-anomalous traffic, or decrease priority of traffic determined anomalous based on anomalies detected by Traffic Anomaly Detector **223**. For example, for traffic associated with a source designated as priority, if the Traffic Anomaly Detector **223** determines that traffic to begin having anomalous characteristics (e.g., a large change in volume or type of traffic at unexpected times, a change in the target object, etc.) the Priority Generator **225** may reduce the priority designation for data from that source (in a system with multiple layers of priority), or give the traffic a non-priority status. In at least some embodiments, priority generator may generate, based on applying artificial intelligence and/or machine learning techniques to the traffic, priorities for processing the traffic. The generated priorities may be used to generate new policies, to modify Client-specified policies, and similar.

(47) FIG. 3 is a process diagram that illustrates a process for scaling-aware traffic request prioritization for a scalable connection-based resource, according to some embodiments. Various aspects of the illustrated process may be performed by one or more components illustrated in FIGS. 1, 2, 7 and/or 8, such as by one or more components of a Request Priority Manager **110**.

(48) At block **302**, scaling-based capacity of a connection-based resource is monitored, by Connection-Based Resource Scale Capacity Monitor **116** of Request Priority Manager **110**. For example, the Monitor **116** may interact with Scalable Connection-Based Resource **130** to monitor the scaling of the Resource **130** or the Monitor **116** may track responses to requests for additional connections and/or responses to requests to scale resources, made by the Request Priority Manager **110** to the Scalable Connection-Based Resource **130**. In at least some embodiments, the Monitor **116** may monitor the scaling-based capacity and store or use information indicative of a scaling state of the Resource **130**. Metrics, based on the monitoring may be published.

(49) At block **304**, a determination is made whether the scaling-based capacity of the Connection-Based Resource **130** is within a threshold limit of a maximum scale capacity of the Resource **130**. For example, the Request Priority Manager **120** may determine, based on comparison of scale information about the capacity of the Resource **130** (e.g., a state tracked and updated by the Scale Capacity Monitor **116**) to a (configurable) threshold limit of maximum capacity (the threshold limit may be a value less than an actual maximum scale allowed or the actual maximum scale value, in

embodiments), or based on a response from the Connection-Based Resource **130** indicating that the Resource **130** cannot scale further, that a maximum capacity of the scalable connection-based resource, or a capacity approaching the maximum capacity has or has not been reached. If the capacity falls within the threshold limit (block **304**, Within) the process returns to monitoring the capacity (block **302**). But if the capacity exceeds the threshold limit (or if a response is received from the Connection-Based Resource **130** indicating that the Resource **130** cannot scale further) a corresponding priority policy is enabled (block **306**), by Traffic Priority Component **114** of Request Priority Manager **110**, for example. In embodiments, the Request Priority Manager **110** may send messages to the Connection Pooler Instances **222A-W** instructing prioritization-based routing. Traffic Priority Component **114** may, for an enabled priority policy, route request traffic in accordance with the enabled priority policy.

(50) At block **308**, the scaling-based capacity of the Connection-Based Resource is monitored, by Resource Scale Monitor **116**, in embodiments. Metrics, based on the monitoring, may be published. At block **310**, a determination is made whether the scaling capacity of the Scalable Connection-Based Resource **130** has fallen below another threshold limit of capacity of the Resource **130** (the thresholds in blocks **304** and **310** may be different, in embodiments). If the capacity remains above the other threshold limit (e.g., if the scale remains at or near maximum scale) the process returns to block **308** (monitoring). Once the capacity falls below the other threshold limit (block **310**, Below) the priority policy is disabled (block **312**) (e.g., by Traffic Priority Component **114**) and request routing returns to normal (normal meaning routing without prioritization resulting in intentionally dropped or delayed requests) such as round-robin or another load-balancing based scheme. For example, the Request Priority Manager **110** may send messages to the Traffic Priority Components **114A-W** of Connection Pooler Instances **222A-W** instructing prioritization-based routing to be disabled.

(51) FIG. **4** is a process diagram that illustrates a request processing process for a connection pooler instance within the context of scaling-aware traffic request prioritization for a scalable connection-based resource, according to some embodiments. In the following description, the process is described within the context of a Connection Pooler Instance **222A**, but may be practiced in other embodiments.

(52) At block **402**, a request is received (e.g., from a Client **102A** and by Connection Pooler Instance **222A**). If a priority policy is not enabled (block **404**, No) the request is routed normally (block **406**) (by Connection Router **112A**, for example). For the case where the priority policy is enabled (block **404**, Yes) the priority policy is applied to route the request (by Traffic Priority Component **114A**). Using FIG. **1** as an example, to apply the priority policy the Request Priority Manager **110** determines, according to the priority policy, which requests **120** received via the respective front-end connections are sent over the backend connections **140** to the Scalable Connection-Based Resource **130**. In the illustrated example in FIG. **4**, if the request is determined to be a priority (block **408**, Priority) the request is processed, normally (block **406**). If the request is determined to not be a priority (block **408**, Not priority) the request is dropped or buffered until processing via priority policy is disabled (block **410**) or for some time-based delay period, etc. The process iterates for additional requests, and across respective ones of the Connection Pooler Instances **222A-W**.

(53) It is contemplated that more complex prioritization schemes may be implemented (e.g., with numerous levels of various priorities, wherein some traffic is transmitted immediately, other traffic is buffered for various different amounts of time, and some traffic is dropped entirely) without departing from the scope of this disclosure. For example, one prioritization scheme may prioritize requests associated with atomic transactions over requests not associated with an atomic transaction. Different levels of priority may be associated with different functionality: a first level of priority may be associated with high-availability routing, a second level may be associated with long-term buffering, a third priority level may be associated with short-term buffering (dropping

the requests from a buffer after some maximum buffer time threshold) and a fourth priority level may specify that the requests be dropped entirely, guaranteeing that higher priority requests are routed.

(54) FIG. 5 is a process diagram that illustrates a process for scaling connection poolers and a connection-based resource, using scaling-aware traffic request prioritization, according to some embodiments. Scaling-based capacity of a Connection-Based Resource **130** and scaling-based capacity of a Connection Pooler **220** are monitored (block **502**), by a same, single component or process or by a Connection-Based Resource Scale Capacity Monitor **116** and a Connection Pool Scale Manager **228**, respectively, in some examples. At block **504**, a determination is made whether a capacity of either exceeds respective threshold limits of respective maximum scale capacity. In embodiments, each of the Connection-Based Resource **130** and the Connection Pool Scale Manager **228** have different threshold limits (e.g., client-configurable via Policy Interface **218**, or determinable via a Request Traffic AI component **221** which may set or adjust or modify the thresholds, based on detected anomalies or based on traffic history).

(55) If neither threshold is exceeded (block **504**, Neither exceeds) the process returns to monitoring at block **502**. If at least one of the thresholds is exceeded (block **504**, At least one exceeds) a corresponding priority policy is enabled (block **506**). For example, each of the Connection-Based Resource **130** and the Connection-Pooler **220** may be associated with different policies, or with a common policy, in various different embodiments. Different policies may specify different values for the thresholds. At block **508**, scaling-based capacity that caused enablement of the priority policy is monitored. For example, if the threshold for the Connection Pooler **220** was exceeded and caused enablement of the priority policy, the scaling-based capacity of the Connection Pooler is monitored, and if the threshold of the Connection-Based Resource **130** was exceeded and caused enablement of the priority policy, the scaling-based capacity of the Connection-Based Resource **130** is monitored.

(56) At block **510** a determination is made whether the capacity that caused enablement of the priority policy has fallen below another threshold limit of scale capacity. If the capacity remains above the other threshold limit (block **510**, Above) the process returns to monitoring at block **508**. Otherwise, for the case where the capacity that caused the enablement has fallen below the other threshold limit (block **510**, Below) the priority policy is disabled (block **512**). For example, instead of priority-based routing, the process may return to normal processing (e.g., round-robin based or whatever routing was being used prior to the priority-based routing, etc.).

(57) In some embodiments, a scaling capacity of the Connection Pooler **220** or the Connection-Based Resource **130** (learned during the monitoring, or in response to a request for a connection for additional resource, for example) may be updated. For example, the system may monitor how much capacity is left to reach a maximum scale and may store that information, as a state of the capacity.

(58) FIG. 6 is a process diagram that illustrates a scaling process for a connection pool within the context of scaling-aware traffic request prioritization for a scalable connection-based resource, according to some embodiments. At least portions of the process may be performed by a Connection Pool Scale Manager **228**, in embodiments. The following description of FIG. 6 describes features that relate to functionality (e.g., provide information for the functionality described in, or otherwise interact with the components that perform the functionality) illustrated in FIGS. 3 and 5, in embodiments.

(59) Request traffic **120** from Clients **102A-Z** is monitored (block **602**) by Connection Pool Scale Manager **228**, for example. The Connection Pool Scale Manager **228** determines whether the Connection Pool needs to be scaled (block **604**), based on the traffic for example. If not (block **604**, No Scaling) the process returns to monitoring at block **602**. If a determination is made to scale-down the Connection Pool (block **604**, Scale-down) the Connection Pool Scale Manager **228** drops a connection (block **606**). If the dropped connection is the last connection of a Connection Pooler

Instance **222**, the corresponding Connection Pooler Instance **222** may be terminated. Scaling limit information for the Connection Pool is updated, by the Connection Pool Scale Manager **228**, for example.

(60) In embodiments, due to the dropped connection/terminated Connection Pooler Instance the Connection Pool and Connection-based Resource **130** may have more room to scale up. The updated scaling limit information may be used by the separate Request Priority Manager processes (e.g., associated with the monitoring in FIGS. **3** and **5**). For example, scaling down may mean that the Connection Pool is below a maximum scale threshold and can disable priority policy-based routing.

(61) If it is decided to Scale-up (block **604**, Scale-Up) the Connection Pool Scale Manager **228** requests an additional connection(s) from the Connection-Based Resource **130** (block **610**). If the request is successful (block **612**, Yes) the scaling limit information is updated, accordingly (block **608**). The updated scaling limit information may be used by the separate Request Priority Manager processes, as noted above. For example, scaling up may mean that a threshold scale limit is met or exceeded and an appropriate priority policy must be enabled.

(62) If the request is unsuccessful (block **612**, No) the Connection Pool Scale Manager **228** can send a message to the Request Priority Manager **110** to enable an appropriate priority policy (block **614**) and/or a message can be sent from the Request Priority Manager **110** or the Connection Pooler **220** to the Connection-Based Resource **130** to scale-up the Connection-Based Resources **130** (block **616**). In some embodiments, a component distinct from the Connection-Based Resource **130** may send an instruction to the Connection-Based Resource **130** instructing the Connection-Based Resource to scale (e.g., in anticipation of an increase in request traffic that the Connection-Based Resource has not learned of yet, and/or in order to have resources to support additional connections **140**) even if the Connection-Based Resource **130** is implemented with a monitor and scaling manager to manage scaling of the Connection-Based Resource **130**.

(63) In embodiments, by adding a connection **140** to the Connection-Based Resource **130**, all current resource instances at the Connection-Based Resource **130** may be at maximum capacity or scale (or close thereto) for the number of connections they can handle and it may look like Request Traffic **120** is still trending to increase, so the Scalable Connection Pooler **220** and/or Connection-Based Resource **130** may start bringing up a new instance (which may take various amounts of time, leading to delayed or dropped traffic requests). In embodiments, the Priority-Aware Scalable Connection Pooler **220** or the Request Priority Manager may send a request or instruction to the Scalable Connection-Based Resource **130**, requesting additional resources, responsive to an increase in Request Traffic **120**, for example

(64) FIG. **7** is a block diagram that illustrates an example service provider network embodiment for scaling-aware traffic request prioritization for a scalable connection-based resource, according to some embodiments. In the non-exhaustive, illustrated embodiment, various functionality described herein is distributed and performed by various Services of the Service Provider Network **710**. Service Provider Network **710** is illustrated with internal and external Clients **102**, Transactional Database Service **712**, Compute Service **714**, Object Storage Service **716**, Logging Monitoring & Reporting Service **718**, Other Compute Service **713**, Other Storage Service **715**, and Other Service **717** connected to external Client(s) **102**, via Network **710**.

(65) Request Priority Manager **110** is illustrated as implemented, at least partially, via Compute Service **714** (e.g., an event-driven, stateless (sometimes referred to as serverless Compute Service) or a web service that provides secure, resizable compute capacity in the cloud, either via virtual, or non-virtual compute instances. In the illustrated embodiment, Connection-Based Resource **130** is illustrated as implemented via Transactional Database Service **712**, but may be a Connection-Based Resource built from one or more other services provided by the Service Provider Network **710**. For example, the Connection-Based Service **130** may be implemented via another type of database service, a managed database service, Relational and Non-Relational Database Services, etc. In

some embodiments the Connection-Based Service **130** may be implemented via a compute service (e.g., an event-driven, stateless (sometimes referred to as serverless Compute Service) or a web service that provides secure, resizable compute capacity in the cloud, either via virtual, or non-virtual compute instances, or may be implemented via a combination of compute and storage services, in some embodiments.

(66) While various external-to-the-service-provider-network Clients **102** and internal-to-the-service-provider-network Clients **102** are illustrated, it is contemplated that other processes, such as micro-services, applications (Client or service provider-based) or the like may be implemented via a combination of compute and storage services of the service provider network, and may act as Clients **102** of the Connection-Based Resource **130**, in various embodiments.

(67) Priority Policy Store **118** is illustrated as implemented at Storage Service **716** (e.g., an object storage service by the service provider, or similar) and Logging, Monitoring & Reporting Service **718** is illustrated as a service provided by the Service Provider Network **710**. In embodiments, the Logging, Monitoring & Reporting Service **718** may perform at least some of the monitoring described herein, and/or may provide metrics for the Request Priority Manager **110** to use in making various determinations described herein.

(68) Example Computer System

(69) FIG. **8** illustrates an example of a computer system, one or more of which may implement various components of scaling-aware traffic request prioritization for a scalable connection-based resource, described and illustrated throughout the disclosure, according to embodiments.

(70) Various portions of systems in FIGS. **1**, **2**, **7** and **8** and/or methods presented in FIGS. **3-6**, described herein, may be executed on one or more computer systems similar to that described herein, which may interact with various other devices of the system.

(71) In the illustrated embodiment, computer system **800** includes one or more processors **810** coupled to a system memory **820** via an input/output (I/O) interface **830**. Computer system **800** further includes a network interface **840** coupled to I/O interface **830**, and one or more input/output devices **860**, such as cursor control device, keyboard, audio device, and display(s). In some embodiments, it is contemplated that embodiments may be implemented using a single instance of computer system **800**, while in other embodiments multiple such systems, or multiple nodes making up computer system **800**, may be configured to host different portions or instances of embodiments. For example, in one embodiment some elements may be implemented via one or more nodes of computer system **800** that are distinct from those nodes implementing other elements.

(72) In various embodiments, computer system **800** may be a uniprocessor system including one processor **810**, or a multiprocessor system including several processors **810** (e.g., two, four, eight, or another suitable number). Processors **810** may be any suitable processor capable of executing instructions. For example, in various embodiments, processors **810** may be general-purpose or embedded processors implementing any of a variety of instruction set architectures (ISAs), such as the x86, PowerPC, SPARC, or MIPS ISAs, or any other suitable ISA. In multiprocessor systems, each of processors **810** may commonly, but not necessarily, implement the same ISA.

(73) In some embodiments, at least one processor **810** may be a graphics processing unit. A graphics processing unit (GPU) may be considered a dedicated graphics-rendering device for a personal computer, workstation, game console or other computer system. GPUs may be very efficient at manipulating and displaying computer graphics and their highly parallel structure may make them more effective than typical CPUs for a range of complex graphical algorithms. For example, a graphics processor may implement a number of graphics primitive operations in a way that makes executing them much faster than drawing directly to the screen with a host central processing unit (CPU). In various embodiments, the methods disclosed herein for scaling-aware traffic request prioritization for a scalable connection-based resource may be implemented by program instructions configured for execution on one of, or parallel execution on two or more of,

such GPUs. The GPU(s) may implement one or more application programmer interfaces (APIs) that permit programmers to invoke the functionality of the GPU(s). Suitable GPUs may be commercially available from vendors such as NVIDIA Corporation, ATI Technologies, and others.

(74) System memory **820** may be configured to store program instructions and/or data accessible by processor **810**. In various embodiments, system memory **820** may be implemented using any suitable memory technology, such as static random-access memory (SRAM), synchronous dynamic RAM (SDRAM), nonvolatile/Flash-type memory, or any other type of memory. In the illustrated embodiment, program instructions and data implementing desired functions, such as those described above for scaling-aware traffic request prioritization for a scalable connection-based resource, are shown stored within system memory **820** as Request Priority Manager **110** and Priority Policy Data Store **118**, respectively. In other embodiments, program instructions and/or data may be received, sent, or stored upon different types of computer-accessible media or on similar media separate from system memory **820** or computer system **800**. Generally speaking, a computer-accessible medium may include storage media or memory media such as magnetic or optical media, e.g., disk or CD/DVD-ROM coupled to computer system **800** via I/O interface **830**. Program instructions and data stored via a computer-accessible medium may be transmitted by transmission media or signals such as electrical, electromagnetic, or digital signals, which may be conveyed via a communication medium such as a network and/or a wireless link, such as may be implemented via network interface **840**. Program instructions may include instructions for implementing the techniques described with respect to any of the FIGS.

(75) In some embodiments, I/O interface **830** may be configured to coordinate I/O traffic between processor **810**, system memory **820**, and any peripheral devices in the device, including network interface **840** or other peripheral interfaces, such as input/output devices. In some embodiments, I/O interface **830** may perform any necessary protocol, timing, or other data transformations to convert data signals from one component (e.g., system memory **820**) into a format suitable for use by another component (e.g., processor **810**). In some embodiments, I/O interface **830** may include support for devices attached through various types of peripheral buses, such as a variant of the Peripheral Component Interconnect (PCI) bus standard or the Universal Serial Bus (USB) standard, for example. In some embodiments, the function of I/O interface **830** may be split into two or more separate components. In addition, in some embodiments some or all of the functionality of I/O interface **830**, such as an interface to system memory **820**, may be incorporated directly into processor **810**.

(76) Network interface **840** may be configured to allow data to be exchanged between computer system **800** and other devices attached to a network **810**, such as other computer systems, or between nodes of computer system **800**. In various embodiments, network interface **840** may support communication via wired or wireless general data networks, such as any suitable type of Ethernet network, for example; via telecommunications/telephony networks such as analog voice networks or digital fiber communications networks; via storage area networks such as Fiber Channel SANs, or via any other suitable type of network and/or protocol.

(77) Computing device **800** may include input/output devices that may, in some embodiments, include one or more display terminals, keyboards, keypads, touchpads, scanning devices, voice or optical recognition devices, accelerometers, multi-touch screens, or any other devices suitable for entering or retrieving data by one or more computer system **800**. Multiple input/output devices may be present in computer system **800** or may be distributed on various nodes of computer system **800**. In some embodiments, similar input/output devices may be separate from computer system **800** and may interact with one or more nodes of computer system **800** through a wired or wireless connection, such as over network interface **840**.

(78) Memory **820** may include program instructions (e.g., such as Request Priority Manager **110**) and data (e.g., Priority Policies in Priority Policy Data Store **118**) accessible by the program instructions. In one embodiment, program instructions may include software elements of a method

illustrated in the above figures. Data storage **118** may include data that may be used in embodiments described herein. In other embodiments, other or different software elements and/or data may be included.

(79) Those skilled in the art will appreciate that computer system **800** is merely illustrative and is not intended to limit the scope of as the systems and methods described herein. In particular, the computer system and devices may include any combination of hardware or software that can perform the indicated functions, including computers, network devices, internet appliances, PDAs, wireless phones, pagers, etc. Computer system **800** may also be connected to other devices that are not illustrated, or instead may operate as a stand-alone system. In addition, the functionality provided by the illustrated components may in some embodiments be combined in fewer components or distributed in additional components. Similarly, in some embodiments, the functionality of some of the illustrated components may not be provided and/or other additional functionality may be available.

(80) Those skilled in the art will also appreciate that, while various items are illustrated as being stored in memory or on storage while being used, these items or portions of them may be transferred between memory and other storage devices for purposes of memory management and data integrity. Alternatively, in other embodiments some or all of the software components may execute in memory on another device and communicate with the illustrated computer system via inter-computer communication. Some or all of the system components or data structures may also be stored (e.g., as instructions or structured data) on a computer-accessible medium or a portable article to be read by an appropriate drive, various examples of which are described above. In some embodiments, instructions stored on a computer-accessible medium separate from computer system **800** may be transmitted to computer system **800** via transmission media or signals such as electrical, electromagnetic, or digital signals, conveyed via a communication medium such as a network and/or a wireless link. Various embodiments may further include receiving, sending, or storing instructions and/or data implemented in accordance with the foregoing description upon a computer-accessible medium. Accordingly, the present solution may be practiced with other computer system configurations. In some embodiments, portions of the techniques described herein (e.g., scaling-aware traffic request prioritization for a scalable connection-based resource) may be hosted in a cloud computing infrastructure.

(81) Various embodiments may further include receiving, sending, or storing instructions and/or data implemented in accordance with the foregoing description upon a computer-accessible medium. Generally speaking, a computer-accessible/readable storage medium may include a non-transitory storage media such as magnetic or optical media, (e.g., disk or DVD/CD-ROM), volatile or non-volatile media such as RAM (e.g., SDRAM, DDR, RDRAM, SRAM, etc.), ROM, etc., as well as transmission media or signals such as electrical, electromagnetic, or digital signals, conveyed via a communication medium such as network and/or a wireless link.

(82) The various methods as illustrated in the figures and described herein represent exemplary embodiments of methods. The methods may be implemented in software, hardware, or a combination thereof. The order of method may be changed, and various elements may be added, reordered, combined, omitted, modified, etc.

(83) Various modifications and changes may be made as would be obvious to a person skilled in the art having the benefit of this disclosure. It is intended to embrace all such modifications and changes and, accordingly, the above description to be regarded in an illustrative rather than a restrictive sense.

Claims

1. A system, comprising: one or more computers comprising respective processors and memory configured to implement a first request priority manager for a scalable connection-based resource,

the first request priority manager configured to: receive request processing traffic for a scalable connection-based resource; monitor a connection capacity or request processing capacity of the scalable connection-based resource; in response to a determination that the scalable connection-based resource is at or within a threshold limit of a maximum capacity for the connection capacity or request processing capacity of the scalable connection-based resource, begin applying a first priority policy obtained from a policy data store to newly received request traffic for the scalable connection-based resource to determine a priority of some of the newly received requests over others; and in response to a determination that the scalable connection-based resource is no longer at or within the threshold limit of the maximum capacity for the scalable connection-based resource, cease applying the first priority policy to newly received request traffic for the scalable connection-based resource.

2. The system of claim 1, wherein the first request priority manager comprises a connection pooler configured to: establish respective front-end connections with a plurality of requesters for receiving the request traffic for the scalable connection-based resource; and maintain a pool of backend connections to the scalable connection-based resource for sending the request traffic to the scalable connection-based resource; wherein to apply the first priority policy the first request priority manager is requests received via the respective front-end connections are sent over the backend connections to the scalable connection-based resource.

3. The system of claim 2, wherein, responsive to an increase in the request traffic for the scalable connection-based resource, the connection pooler is configured to request one or more additional backend connections to the scalable connection-based resource to increase a number of connections in the pool of backend connections, wherein the determination that the scalable connection-based resource is at or within the threshold limit of the maximum capacity is based on the ability of the scalable connection-based resource to accept the one or more additional backend connections.

4. The system of claim 3, wherein the scalable connection-based resource comprises a group of resources instances configured to establish the backend connections with the connection pooler and process requests from the request traffic received via respective connections of the backend connections, wherein the scalable connection-based resource is scalable to increase or decrease a number of resources instances in the group of resources instances; and wherein the first request priority manager is configured to request the scalable connection-based resource to increase the number of resources instances in the group of resources instances in response to an indication that the scalable connection-based resource is unable to accept the one or more additional backend connections.

5. The system of claim 2, further comprising: one or more computers comprising respective processors and memory configured to implement a second request priority manager, implemented at the scalable connection-based resource and configured to: receive request traffic for the scalable connection-based resource over the backend connections; monitor the connection capacity or request processing capacity of the scalable connection-based resource; and in response to a determination that the scalable connection-based resource is at or within a threshold limit of a maximum capacity for the scalable connection-based resource, begin applying a second priority policy to newly received request traffic for the scalable connection-based resource to prioritize some requests over others, wherein the second priority policy is different from the first priority policy.

6. A method, performed by one or more processors of one or more computing devices, the method comprising: receiving, by a first request priority manager, request traffic for a scalable connection-based resource; monitoring a connection capacity or request processing capacity of the scalable connection-based resource; applying, by the first request priority manager and in response to a determination that the scalable connection-based resource is at or within a threshold limit of a maximum capacity for the connection capacity or request processing capacity of the scalable connection-based resource, a first priority policy obtained from a policy data store to newly

received request traffic for the scalable connection-based resource to determine a priority of some of the newly received requests over others; and in response to determining that the scalable connection-based resource is no longer at or within the threshold limit of the maximum capacity for the scalable connection-based resource, cease applying the first priority policy to newly received request traffic for the scalable connection-based resource.

7. The method of claim 6, wherein the first request priority manager comprises a connection pooler configured to: establishing, by a connection pooler of the first request priority manager, respective front-end connections with a plurality of requesters for receiving the request traffic for the scalable connection-based resource; and maintaining, by the connection pooler, a pool of backend connections to the scalable connection-based resource for sending the request traffic to the scalable connection-based resource; wherein said applying the first priority policy comprises determining, according to the first priority policy, which requests received via the respective front-end connections are sent over the backend connections to the scalable connection-based resource.

8. The method of claim 7, further comprising: responsive to an increase in the request traffic for the scalable connection-based resource, requesting by the connection pooler one or more additional backend connections to the scalable connection-based resource to increase a number of connections in the pool of backend connections, wherein said determining that the scalable connection-based resource is at or within the threshold limit of the maximum capacity is based on the ability of the scalable connection-based resource to accept the one or more additional backend connections.

9. The method of claim 8, wherein: the connection pooler comprises a group of connection pool instances for maintaining respective connections of the pool of backend connections to the scalable connection-based resource; said increase the number of connections in the pool of backend connections comprises increase a number of connection pool instances in the group of connection pool instances; and the method further comprises: begin applying, by the first request priority manager, the first priority policy to the request traffic responsive to reaching a maximum number of connection pool instances or a maximum number of connections in the pool of backend connections.

10. The method of claim 6, wherein: the first priority policy specifies one or more data objects stores at the scalable connection-based resource; and said applying the first priority policy comprises prioritizing requests that access at least one of the one or more data objects specified in the first priority policy.

11. The method of claim 6, wherein: the first priority policy specifies one or more requestor identifiers (IDs); and said applying the first priority policy comprises prioritizing respective requests that indicate one of the one or more requestor IDs specified in the first priority policy as a source of the respective request.

12. The method of claim 6, further comprising: providing, by the first request priority manager, an interface for a user of the scalable connection-based resource to specify the first priority policy.

13. The method of claim 6, wherein: the first request priority manager comprises an artificial intelligence component; and the method further comprises specifying or modifying, by the artificial intelligence component, the first priority policy based on historical request traffic to the scalable connection-based resource.

14. One or more non-transitory computer-readable media storing program instructions executable on or across one or more processors implement a first request priority manager to: monitor a connection capacity or request processing capacity of a scalable connection-based resource to which request traffic is directed; begin application, in response to a determination that the scalable connection-based resource is at or within a threshold limit of a maximum capacity for the connection capacity or request processing capacity of the scalable connection-based resource, of a first priority policy obtained from a policy data store to newly received request traffic for the scalable connection-based resource to determine a priority of some of the newly received requests over others; and cease application, in response to a determination that the scalable connection-

based resource is no longer at or within the threshold limit of the maximum capacity for the scalable connection-based resource, of the first priority policy to newly received request traffic for the scalable connection-based resource.

15. The one or more non-transitory computer-readable media of claim 14, storing program instructions executable to implement a connection pooler at the first request priority manager, the connection pooler to: establish respective front-end connections with a plurality of requesters for receiving the request traffic for the scalable connection-based resource; and maintain a pool of backend connections to the scalable connection-based resource for sending the request traffic to the scalable connection-based resource; wherein said application of the first priority policy comprises determine, according to the first priority policy, which requests received via the respective front-end connections are sent over the backend connections to the scalable connection-based resource, and which requests received via the respective front-end connections are not sent over the backend connections to the scalable connection-based resource.

16. The one or more non-transitory computer-readable media of claim 15, storing program instructions executable to cause the connection pooler to: responsive to an increase in the request traffic for the scalable connection-based resource, request one or more additional backend connections to the scalable connection-based resource to increase a number of connections in the pool of backend connections; and base the determination that the scalable connection-based resource is at or within the threshold limit of the maximum capacity on the ability of the scalable connection-based resource to accept the one or more additional backend connections.

17. The one or more non-transitory computer-readable media of claim 15, wherein: the connection pooler comprises a group of connection pool instances for maintaining respective connections of the pool of backend connections to the scalable connection-based resource; to increase the number of connections in the pool of backend connections, the program instructions are executable to increase a number of connection pool instances in a group of connection pool instances; and the program instructions are executable to perform said application, by the first request priority manager, of the first priority policy to the request traffic responsive to reaching a maximum number of connection pool instances or a maximum number of connections in the pool of backend connections.

18. The one or more non-transitory computer-readable media of claim 15, wherein: the first priority policy specifies one or more data objects stores at the scalable connection-based resource; and to perform said application of the first priority policy the program instructions are executable to prioritize requests that access at least one of the one or more data objects specified in the first priority policy.

19. The one or more non-transitory computer-readable media of claim 15, wherein: the first priority policy specifies one or more requestor identifiers (IDs); and to perform said application of the first priority policy the program instructions are executable to prioritize respective requests that indicate one of the one or more requestor IDs specified in the first priority policy as a source of the respective request.

20. The one or more non-transitory computer-readable media of claim 15, wherein the program instructions are executable to cause the first request priority manager to determine an anomalous pattern in the request traffic and modify the first priority policy to change a prioritization for new request traffic corresponding to the determined anomalous pattern.
