

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12395531
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Deshmukh; Pooja

Cloud access security broker user interface and analytics systems and methods

Abstract

Systems and methods include, providing a UI for a tenant to input one or more malware and DLP rules, and trusted user exceptions; responsive to a scan by the CASB system of a plurality of users associated with a tenant in a SaaS application where the scan includes identifying malware in content in the SaaS application and performing DLP in the content in the SaaS application based on the one or more malware and DLP rules and trusted user exceptions, maintaining records associated with a plurality of incidents for the malware and the DLP; and providing the UI for the tenant including an analytics view with a plurality of summary tiles including visualizations of the plurality of incidents for the malware and DLP for the tenant and a table listing any of the plurality of incidents for the malware and the DLP for the tenant.

Inventors:	Deshmukh; Pooja (Sunnyvale, CA)
Applicant:	Zscaler, Inc. (San Jose, CA)
Family ID:	1000008767505
Assignee:	Zscaler, Inc. (San Jose, CA)
Appl. No.:	17/949522
Filed:	September 21, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20230018809 A1	Jan. 19, 2023

Related U.S. Application Data

continuation-in-part parent-doc US 16950136 20201117 US 12041089 child-doc US 17949522

Publication Classification

Int. Cl.: H04L9/40 (20220101)

U.S. Cl.:

CPC H04L63/20 (20130101); H04L63/1433 (20130101); H04L63/145 (20130101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
8185510	12/2011	Chaudhry et al.	N/A	N/A
8429111	12/2012	Kailash et al.	N/A	N/A
8869259	12/2013	Udupa et al.	N/A	N/A
8887249	12/2013	Schekochikhin et al.	N/A	N/A
9152789	12/2014	Natarajan et al.	N/A	N/A
9154475	12/2014	Kailash et al.	N/A	N/A
9473537	12/2015	Sinha et al.	N/A	N/A
9609015	12/2016	Natarajan et al.	N/A	N/A
10523710	12/2018	Sinha et al.	N/A	N/A
2017/0214701	12/2016	Hasan	N/A	H04L 63/1433
2017/0372070	12/2016	Burdett	N/A	G06F 21/567
2020/0259792	12/2019	Devarajan	N/A	G06F 21/554
2020/0389496	12/2019	Xuan	N/A	H04L 63/10

OTHER PUBLICATIONS

Tahboub et al., “Data Leakage/Loss Prevention Systems (DLP)”, Oct. 24, 2014, IEEE, 2014 World Congress on Computer Applications and Information Systems (WCCAIS) (2014, pp. 1-6) (Year: 2014). cited by examiner

Primary Examiner: Zand; Davoud A

Attorney, Agent or Firm: Baratta Law PLLC

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION(S) (1) The present application is a continuation-in-part of U.S. patent application Ser. No. 16/950,136, filed Nov. 17, 2020, and entitled “CLOUD ACCESS SECURITY BROKER USER INTERFACE SYSTEMS AND METHODS” the contents of which are incorporated by reference.

FIELD OF THE DISCLOSURE

(1) The present disclosure relates generally to networking and computing. More particularly, the

present disclosure relates to Cloud Access Security Broker (CASB) User Interface (UI) systems and methods.

BACKGROUND OF THE DISCLOSURE

(2) Traditionally, before the cloud, corporate or enterprise resources were fully under the control of Information Technology (IT) administration (“admins”). That is, sensitive enterprise data was located within a network under IT admin control with perimeter defenses. Here, IT admins have full control of access privileges, activity, etc. As is well-known, enterprises are moving their IT infrastructure to the cloud for a variety of cloud services (Software-as-a-Service (SaaS)) for email (e.g., Office 365, Gmail, etc.), file storage (OneDrive, Dropbox, Box, Google Drive, SharePoint, etc.), document preparation and content collaboration (e.g., Office 365, Google Docs, etc.), Customer Relationship Management (CRM) (e.g., Salesforce, etc.), and the like. Here, enterprise IT admins no longer have the same level of control of enterprise resources, i.e., content is stored in the cloud, and IT simply does not have the same level of control as before.

(3) A Cloud Access Security Broker (CASB) is an on-premises system or cloud-based service between cloud service users and cloud applications. The CASB is configured to monitor activity and enforce security policies, such as monitoring user activity, warning administrators about potentially hazardous actions, Data Loss Prevention (DLP), enforcing security policy compliance, automatically preventing malware, etc. For example, a CASB system, either on-premises or as a cloud-based service, can scan through a large number of files in a cloud or SaaS application, e.g., Office 365, Dropbox, Box, Google Drive, Salesforce, etc. This places tremendous loads on traditional CASB systems, resulting in latency, inability to properly scan all files, poor user experience, etc. In effect, an objective of a CASB system or scanner is to provide IT admin with control as if the enterprise resources were fully under the IT admin's control as before the cloud.

(4) Again, the goal of IT is to have similar control of cloud-based deployments as was with conventional deployments. Towards this goal, there is a need to support visualizations, reporting, a User Interface (UI), etc. to support investigation and remediation. With the existing solutions, CASBs are able to provide visibility into cloud usage throughout the organization, control access to cloud services, and threat prevention. However, with the incident-based reports, there lacks a concrete view of the top objects that threaten the organization's security.

BRIEF SUMMARY OF THE DISCLOSURE

(5) The present disclosure relates to Cloud Access Security Broker (CASB) User Interface (UI) systems and methods. The CASB UI systems and methods include consolidated organization level insights into organization's sensitive data exposed accidentally or intentionally outside of the acceptable business sharing policies. The three primary entities analyzed in this report are—unique data objects, unique users internal to the organization and unique external entities. The report presented herein provide insights into potential compliance and user behavioral issues at a click of the button. Efficiently, practitioners can find out if their organization's data is externally exposed (Files Views) or if external malicious entities have access to it (Collaborators Views). The easy to consume summary tiles highlight the severity of the problem at hand and practitioners can investigate further details using the incident reporting drill-down mechanism. It exposes potential data sharing compliance issues in an easily actionable manner. It also highlights potential user behavior issues (Owners Views) where employees repeatedly share data outside of formal business sharing policies. It helps organizations to identify gaps in user education and provide targeted mitigation plans. The history of such repeated behavior also helps organizations in identifying potential malicious insiders. In addition to the incidents generated from cloud applications in use, the report shows which unique objects, internal/external users are associated with DLP violations or malware, so that they can investigate and take actions accordingly.

(6) In various embodiments, the present disclosure includes a method with steps, a system including a cloud node, a Cloud Access Security Broker (CASB) system, and/or a cloud-based system configured to implement the steps, and a non-transitory computer-readable storage medium

having computer-readable code stored thereon for programming one or more processors to perform the steps. The steps include, providing a User Interface (UI) for a tenant to input one or more malware and Data Loss Prevention (DLP) rules, and trusted user exceptions, wherein the trusted user exceptions identify one or more specific users and rule exceptions for the specific users; responsive to a scan by the CASB system of a plurality of users associated with a tenant in a Software-as-a-Service (SaaS) application where the scan includes identifying malware in content in the SaaS application and performing DLP in the content in the SaaS application based on the one or more malware and DLP rules and trusted user exceptions, maintaining records associated with a plurality of incidents for the malware and the DLP; and providing the UI for the tenant including an analytics view with a plurality of summary tiles including visualizations of the plurality of incidents for the malware and DLP for the tenant and a table listing any of the plurality of incidents for the malware and the DLP for the tenant, including any of unique data objects, unique users internal to the tenant, and unique external entities, associated with the plurality of incidents.

(7) The steps can further include providing the UI for the tenant to onboard a plurality of SaaS applications including the SaaS application. The steps can further include providing the UI for the tenant to configure policies for the DLP and for the malware for the SaaS application. The steps can further include, responsive to a selection of any entry in the table, providing a popup listing details associated with the corresponding incident. The SaaS application can be one of a plurality of SaaS applications for the tenant, and wherein the visualizations can include a table listing the plurality of incidents associated with the plurality of SaaS applications. The visualizations can include one or more pie charts illustrating the plurality of incidents. The visualizations can include a line chart illustrating the plurality of incidents over time. The steps can further include performing remediation of the plurality of incidents, wherein the remediation of the plurality of incidents includes granular remediation, or wherein the remediation of the plurality of incidents includes tombstoning one or more files for better user experience. A scheduler can be configured to control historic and ongoing scans of content of the plurality of users.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The present disclosure is illustrated and described herein with reference to the various drawings, in which like reference numbers are used to denote like system components/method steps, as appropriate, and in which:
- (2) FIG. 1A is a network diagram of a cloud-based system offering security as a service;
- (3) FIG. 1B is a network diagram of an example implementation of the cloud-based system of FIG. 1;
- (4) FIG. 2 is a block diagram of a server which may be used in the cloud-based system of FIG. 1, to implement a CASB system or the like;
- (5) FIG. 3 is a block diagram of a mobile device which may be used in the cloud-based system of FIG. 1 or the like;
- (6) FIG. 4 is a network diagram of a CASB system;
- (7) FIG. 5 is a functional block diagram of filing crawling of the SaaS provider with the CASB system;
- (8) FIG. 6 is a flowchart of a file crawling process based on a change log;
- (9) FIG. 7 is a flowchart of a file crawling process based on breadth-first traversal;
- (10) FIG. 8 is a flow diagram of example operations between the CASB client, the controller, the message broker, a worker, and the SaaS provider;
- (11) FIG. 9 is a flow diagram of an architecture of a CASB-webhooks system;
- (12) FIG. 10 is a flowchart of a CASB-webhooks integration process, that may be implemented

- through the CASB-webhooks system of FIG. 10, or in other approaches;
- (13) FIG. 11 is a flow diagram of subscription and renewal process for the registration step and the renewal step in the CASB-webhooks integration process;
- (14) FIG. 12 is a map illustrating an example system including two CASB systems geographically distributed and two SaaS applications also geographically distributed;
- (15) FIG. 13 is a flowchart of a historical and live scanning process for CASB functionality;
- (16) FIG. 14 is a block diagram of a CASB in-memory data store system;
- (17) FIG. 15 is a flowchart of a record processing process implemented in the data store of the CASB in-memory data store system of FIG. 14;
- (18) FIG. 16 is a diagram of an example implementation of the filehash from the data store of the CASB in-memory data store system of FIG. 14;
- (19) FIG. 17 is a diagram of data duplication prevention in the data store;
- (20) FIGS. 18-23 are screenshots of an onboarding process for enrolling an organization through the CASB system with a SaaS application;
- (21) FIGS. 24-27 are screenshots of a policy configuration process, subsequent to the onboarding process;
- (22) FIGS. 28-35 are screenshots of various reports and visualizations associated with monitoring through the CASB system;
- (23) FIG. 36 is a screenshot of an assets with incidents view illustrating a file tab;
- (24) FIG. 37 is a screenshot of the assets with incidents view illustrating an owners tab;
- (25) FIG. 38 is a screenshot of a popup view illustrating files with incidents for a specific owner;
- (26) FIG. 39 is a screenshot of the assets with incidents view illustrating a collaborators tab;
- (27) FIG. 40 is a screenshot of a popup for configuring one or more trusted domains or trusted users;
- (28) FIG. 41 is a screenshot of a dropdown list of options for selecting trusted domains;
- (29) FIG. 42 is a screenshot of a security exception input process;
- (30) FIG. 43 is a screenshot of a list of tenants representing the number of trusted domains and trusted users for each tenant;
- (31) FIG. 44 is a flowchart of a CASB UI process; and
- (32) FIG. 45 is a flowchart of a CASB UI process.

DETAILED DESCRIPTION OF THE DISCLOSURE

(33) Again, the present disclosure relates to Cloud Access Security Broker (CASB) User Interface (UI) systems and methods. The CASB UI systems and methods include consolidated organization level insights into organization's sensitive data exposed accidentally or intentionally outside of the acceptable business sharing policies. The three primary entities analyzed in this report are—unique data objects, unique users internal to the organization and unique external entities. The report presented herein provide insights into potential compliance and user behavioral issues at a click of the button. Efficiently, practitioners can find out if their organization's data is externally exposed (Files Views and Email Messages Views) or if external malicious entities have access to it (Collaborators Views). For example, the organization's data can be objects which may include a file, such as in the case of a file sharing SaaS, an email message in the case of email SaaS applications (e.g., FIG. 33 are email message based object views), etc. The easy to consume summary tiles highlight the severity of the problem at hand and practitioners can investigate further details using the incident reporting drill-down mechanism. It exposes potential data sharing compliance issues in an easily actionable manner. It also highlights potential user behavior issues (Owners Views) where employees repeatedly share data outside of formal business sharing policies. It helps organizations to identify gaps in user education and provide targeted mitigation plans. The history of such repeated behavior also helps organizations in identifying potential malicious insiders. In addition to the incidents generated from cloud applications in use, the report shows which unique objects, internal/external users are associated with DLP violations or malware,

so that they can investigate and take actions accordingly.

(34) Example Cloud-Based System Architecture

(35) FIG. 1A is a network diagram of a cloud-based system **100** offering security as a service. Specifically, the cloud-based system **100** can offer a Secure Internet and Web Gateway as a service to various users **102**, as well as other cloud services. In this manner, the cloud-based system **100** is located between the users **102** and the Internet as well as any cloud services **106** (or applications) accessed by the users **102**. As such, the cloud-based system **100** provides inline monitoring inspecting traffic between the users **102**, the Internet **104**, and the cloud services **106**, including Secure Sockets Layer (SSL) traffic. The cloud-based system **100** can offer access control, threat prevention, data protection, etc. The access control can include a cloud-based firewall, cloud-based intrusion detection, Uniform Resource Locator (URL) filtering, bandwidth control, Domain Name System (DNS) filtering, etc. The threat prevention can include cloud-based intrusion prevention, protection against advanced threats (malware, spam, Cross-Site Scripting (XSS), phishing, etc.), cloud-based sandbox, antivirus, DNS security, etc. The data protection can include Data Loss Prevention (DLP), cloud application security such as via Cloud Access Security Broker (CASB), file type control, etc.

(36) The cloud-based firewall can provide Deep Packet Inspection (DPI) and access controls across various ports and protocols as well as being application and user aware. The URL filtering can block, allow, or limit website access based on policy for a user, group of users, or entire organization, including specific destinations or categories of URLs (e.g., gambling, social media, etc.). The bandwidth control can enforce bandwidth policies and prioritize critical applications such as relative to recreational traffic. DNS filtering can control and block DNS requests against known and malicious destinations.

(37) The cloud-based intrusion prevention and advanced threat protection can deliver full threat protection against malicious content such as browser exploits, scripts, identified botnets and malware callbacks, etc. The cloud-based sandbox can block zero-day exploits (just identified) by analyzing unknown files for malicious behavior. Advantageously, the cloud-based system **100** is multi-tenant and can service a large volume of the users **102**. As such, newly discovered threats can be promulgated throughout the cloud-based system **100** for all tenants practically instantaneously. The antivirus protection can include antivirus, antispware, antimalware, etc. protection for the users **102**, using signatures sourced and constantly updated. The DNS security can identify and route command-and-control connections to threat detection engines for full content inspection.

(38) The DLP can use standard and/or custom dictionaries to continuously monitor the users **102**, including compressed and/or SSL-encrypted traffic. Again, being in a cloud implementation, the cloud-based system **100** can scale this monitoring with near-zero latency on the users **102**. The cloud application security can include CASB functionality to discover and control user access to known and unknown cloud services **106**. The file type controls enable true file type control by the user, location, destination, etc. to determine which files are allowed or not. A description of DLP functionality is provided in commonly-assigned U.S. patent application Ser. No. 16/923,225, filed Jul. 8, 2020, and entitled "Data Loss Prevention via Indexed Document Management," the contents of which are incorporated by reference herein in their entirety.

(39) For illustration purposes, the users **102** of the cloud-based system **100** can include a mobile device **110**, a headquarters (HQ) **112** which can include or connect to a data center (DC) **114**, Internet of Things (IoT) devices **116**, a branch office/remote location **118**, etc., and each includes one or more user devices (an example user device **300** is illustrated in FIG. 3). The devices **110**, **116**, and the locations **112**, **114**, **118** are shown for illustrative purposes, and those skilled in the art will recognize there are various access scenarios and other users **102** for the cloud-based system **100**, all of which are contemplated herein. The users **102** can be associated with a tenant, which may include an enterprise, a corporation, an organization, etc. That is, a tenant or company is a group of users who share a common access with specific privileges to the cloud-based system **100**,

a cloud service, etc. In an embodiment, the headquarters **112** can include an enterprise's network with resources in the data center **114**. The mobile device **110** can be a so-called road warrior, i.e., users that are off-site, on-the-road, etc. Further, the cloud-based system **100** can be multi-tenant, with each tenant having its own users **102** and configuration, policy, rules, etc. One advantage of the multi-tenancy and a large volume of users is the zero-day/zero-hour protection in that a new vulnerability can be detected and then instantly remediated across the entire cloud-based system **100**. The same applies to policy, rule, configuration, etc. changes—they are instantly remediated across the entire cloud-based system **100**. As well, new features in the cloud-based system **100** can also be rolled up simultaneously across the user base, as opposed to selective and time-consuming upgrades on every device at the locations **112**, **114**, **118**, and the devices **110**, **116**.

(40) Logically, the cloud-based system **100** can be viewed as an overlay network between users (at the locations **112**, **114**, **118**, and the devices **110**, **106**) and the Internet **104** and the cloud services **106**. Previously, the IT deployment model included enterprise resources and applications stored within the data center **114** (i.e., physical devices) behind a firewall (perimeter), accessible by employees, partners, contractors, etc. on-site or remote via Virtual Private Networks (VPNs), etc. The cloud-based system **100** is replacing the conventional deployment model. The cloud-based system **100** can be used to implement these services in the cloud without requiring the physical devices and management thereof by enterprise IT administrators. As an ever-present overlay network, the cloud-based system **100** can provide the same functions as the physical devices and/or appliances regardless of geography or location of the users **102**, as well as independent of platform, operating system, network access technique, network access provider, etc.

(41) There are various techniques to forward traffic between the users **102** at the locations **112**, **114**, **118**, and via the devices **110**, **116**, and the cloud-based system **100**. Typically, the locations **112**, **114**, **118** can use tunneling where all traffic is forward through the cloud-based system **100**. For example, various tunneling protocols are contemplated, such as Generic Routing Encapsulation (GRE), Layer Two Tunneling Protocol (L2TP), Internet Protocol (IP) Security (IPsec), customized tunneling protocols, etc. The devices **110**, **116** can use a local application that forwards traffic, a proxy such as via a Proxy Auto-Config (PAC) file, and the like. A key aspect of the cloud-based system **100** is all traffic between the users **102** and the Internet **104** or the cloud services **106** is via the cloud-based system **100**. As such, the cloud-based system **100** has visibility to enable various functions, all of which are performed off the user device in the cloud.

(42) The cloud-based system **100** can also include a management system **120** for tenant access to provide global policy and configuration as well as real-time analytics. This enables IT administrators to have a unified view of user activity, threat intelligence, application usage, etc. For example, IT administrators can drill-down to a per-user level to understand events and correlate threats, to identify compromised devices, to have application reility, and the like. The cloud-based system **100** can further include connectivity to an Identity Provider (IDP) **122** for authentication of the users **102** and to a Security Information and Event Management (SIEM) system **124** for event logging. The system **124** can provide alert and activity logs on a per-user **102** basis.

(43) FIG. 1B is a network diagram of an example implementation of the cloud-based system **100**. In an embodiment, the cloud-based system **100** includes a plurality of enforcement nodes (EN) **150**, labeled as enforcement nodes **150-1**, **150-2**, **150-N**, interconnected to one another and interconnected to a central authority (CA) **152**. The nodes **150**, **152**, while described as nodes, can include one or more servers, including physical servers, virtual machines (VM) executed on physical hardware, etc. That is, a single node **150**, **152** can be a cluster of devices. An example of a server is illustrated in FIG. 1B. The cloud-based system **100** further includes a log router **154** that connects to a storage cluster **156** for supporting log maintenance from the enforcement nodes **150**. The central authority **152** provide centralized policy, real-time threat updates, etc. and coordinates the distribution of this data between the enforcement nodes **150**. The enforcement nodes **150** provide an onramp to the users **102** and are configured to execute policy, based on the central

authority **152**, for each user **102**. The enforcement nodes **150** can be geographically distributed, and the policy for each user **102** follows that user **102** as he or she connects to the nearest (or other criteria) enforcement node **150**.

(44) The enforcement nodes **150** are full-featured secure internet gateways that provide integrated internet security. They inspect all web traffic bi-directionally for malware and enforce security, compliance, and firewall policies, as described herein. In an embodiment, each enforcement node **150** has two main modules for inspecting traffic and applying policies: a web module and a firewall module. The enforcement nodes **150** are deployed around the world and can handle hundreds of thousands of concurrent users with millions of concurrent sessions. Because of this, regardless of where the users **102** are, they can access the Internet **104** from any device, and the enforcement nodes **150** protect the traffic and apply corporate policies. The enforcement nodes **150** can implement various inspection engines therein, and optionally, send sandboxing to another system. The enforcement nodes **150** include significant fault tolerance capabilities, such as deployment in active-active mode to ensure availability and redundancy as well as continuous monitoring.

(45) In an embodiment, customer traffic is not passed to any other component within the cloud-based system **100**, and the enforcement nodes **150** can be configured never to store any data to disk. Packet data is held in memory for inspection and then, based on policy, is either forwarded or dropped. Log data generated for every transaction is compressed, tokenized, and exported over secure TLS connections to the log routers **154** that direct the logs to the storage cluster **156**, hosted in the appropriate geographical region, for each organization.

(46) The central authority **152** hosts all customer (tenant) policy and configuration settings. It monitors the cloud and provides a central location for software and database updates and threat intelligence. Given the multi-tenant architecture, the central authority **152** is redundant and backed up in multiple different data centers. The enforcement nodes **150** establish persistent connections to the central authority **152** to download all policy configurations. When a new user connects to an enforcement node **150**, a policy request is sent to the central authority **152** through this connection. The central authority **152** then calculates the policies that apply to that user **102** and sends the policy to the enforcement node **150** as a highly compressed bitmap.

(47) Once downloaded, a tenant's policy is cached until a policy change is made in the management system **120**. When this happens, all of the cached policies are purged, and the enforcement nodes **150** request the new policy when the user **102** next makes a request. In an embodiment, the enforcement node **150** exchange “heartbeats” periodically, so all enforcement nodes **150** are informed when there is a policy change. Any enforcement node **150** can then pull the change in policy when it sees a new request.

(48) The cloud-based system **100** can be a private cloud, a public cloud, a combination of a private cloud and a public cloud (hybrid cloud), or the like. Cloud computing systems and methods abstract away physical servers, storage, networking, etc., and instead offer these as on-demand and elastic resources. The National Institute of Standards and Technology (NIST) provides a concise and specific definition which states cloud computing is a model for enabling convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction. Cloud computing differs from the classic client-server model by providing applications from a server that are executed and managed by a client's web browser or the like, with no installed client version of an application required. Centralization gives cloud service providers complete control over the versions of the browser-based and other applications provided to clients, which removes the need for version upgrades or license management on individual client computing devices. The phrase “Software as a Service” (SaaS) is sometimes used to describe application programs offered through cloud computing. A common shorthand for a provided cloud computing service (or even an aggregation of all existing cloud services) is “the cloud.” The cloud-based system **100** is illustrated herein as an example

embodiment of a cloud-based system, and other implementations are also contemplated.

(49) As described herein, the terms cloud services and cloud applications may be used interchangeably. The cloud service **106** is any service made available to users on-demand via the Internet, as opposed to being provided from a company's on-premises servers. A cloud application, or cloud app, is a software program where cloud-based and local components work together. The cloud-based system **100** can be utilized to provide example cloud services, including Zscaler Internet Access (ZIA), Zscaler Private Access (ZPA), and Zscaler Digital Experience (ZDX), all from Zscaler, Inc. (the assignee and applicant of the present application). The ZIA service can provide the access control, threat prevention, and data protection described above with reference to the cloud-based system **100**. ZPA can include access control, microservice segmentation, etc. The ZDX service can provide monitoring of user experience, e.g., Quality of Experience (QoE), Quality of Service (QoS), etc., in a manner that can gain insights based on continuous, inline monitoring. For example, the ZIA service can provide a user with Internet Access, and the ZPA service can provide a user with access to enterprise resources instead of traditional Virtual Private Networks (VPNs), namely ZPA provides Zero Trust Network Access (ZTNA). Those of ordinary skill in the art will recognize various other types of cloud services **106** are also contemplated. Also, other types of cloud architectures are also contemplated, with the cloud-based system **100** presented for illustration purposes.

(50) Other cloud services can include Office 365, Dropbox, Box, Google Drive, Salesforce, and the like. In the context of these services, a provider of such cloud services can be referred to as a cloud provider, a SaaS provider, etc., and may utilize a hardware architecture similar to the cloud-based system **100**. Of course, other types of cloud architectures are also contemplated.

(51) Example Server Architecture

(52) FIG. 2 is a block diagram of a server **200**, which may be used in the cloud-based system **100**, in a CASB system, in other systems, or standalone. For example, the enforcement nodes **150** and the central authority **152** may be formed as one or more of the servers **200**. The server **200** may be a digital computer that, in terms of hardware architecture, generally includes a processor **202**, input/output (I/O) interfaces **204**, a network interface **206**, a data store **208**, and memory **210**. It should be appreciated by those of ordinary skill in the art that FIG. 3 depicts the server **200** in an oversimplified manner, and a practical embodiment may include additional components and suitably configured processing logic to support known or conventional operating features that are not described in detail herein. The components (**202**, **204**, **206**, **208**, and **210**) are communicatively coupled via a local interface **212**. The local interface **212** may be, for example, but not limited to, one or more buses or other wired or wireless connections, as is known in the art. The local interface **212** may have additional elements, which are omitted for simplicity, such as controllers, buffers (caches), drivers, repeaters, and receivers, among many others, to enable communications. Further, the local interface **212** may include address, control, and/or data connections to enable appropriate communications among the aforementioned components.

(53) The processor **202** is a hardware device for executing software instructions. The processor **202** may be any custom made or commercially available processor, a Central Processing Unit (CPU), an auxiliary processor among several processors associated with the server **200**, a semiconductor-based microprocessor (in the form of a microchip or chipset), or generally any device for executing software instructions. When the server **200** is in operation, the processor **202** is configured to execute software stored within the memory **210**, to communicate data to and from the memory **210**, and to generally control operations of the server **200** pursuant to the software instructions. The I/O interfaces **204** may be used to receive user input from and/or for providing system output to one or more devices or components.

(54) The network interface **206** may be used to enable the server **200** to communicate on a network, such as the Internet **104**. The network interface **206** may include, for example, an Ethernet card or adapter (e.g., 10BaseT, Fast Ethernet, Gigabit Ethernet, 10 GbE) or a Wireless Local Area Network

(WLAN) card or adapter (e.g., 802.11 a/b/g/n/ac). The network interface **206** may include address, control, and/or data connections to enable appropriate communications on the network. A data store **208** may be used to store data. The data store **208** may include any of volatile memory elements (e.g., random access memory (RAM, such as DRAM, SRAM, SDRAM, and the like)), nonvolatile memory elements (e.g., ROM, hard drive, tape, CDROM, and the like), and combinations thereof. Moreover, the data store **208** may incorporate electronic, magnetic, optical, and/or other types of storage media. In one example, the data store **208** may be located internal to the server **200**, such as, for example, an internal hard drive connected to the local interface **212** in the server **200**.

Additionally, in another embodiment, the data store **208** may be located external to the server **200** such as, for example, an external hard drive connected to the I/O interfaces **204** (e.g., SCSI or USB connection). In a further embodiment, the data store **208** may be connected to the server **200** through a network, such as, for example, a network-attached file server.

(55) The memory **210** may include any of volatile memory elements (e.g., random access memory (RAM, such as DRAM, SRAM, SDRAM, etc.)), nonvolatile memory elements (e.g., ROM, hard drive, tape, CDROM, etc.), and combinations thereof. Moreover, the memory **210** may incorporate electronic, magnetic, optical, and/or other types of storage media. Note that the memory **210** may have a distributed architecture, where various components are situated remotely from one another but can be accessed by the processor **202**. The software in memory **210** may include one or more software programs, each of which includes an ordered listing of executable instructions for implementing logical functions. The software in the memory **210** includes a suitable Operating System (O/S) **214** and one or more programs **216**. The operating system **214** essentially controls the execution of other computer programs, such as the one or more programs **216**, and provides scheduling, input-output control, file and data management, memory management, and communication control and related services. The one or more programs **216** may be configured to implement the various processes, algorithms, methods, techniques, etc. described herein.

(56) It will be appreciated that some embodiments described herein may include one or more generic or specialized processors (“one or more processors”) such as microprocessors; Central Processing Units (CPUs); Digital Signal Processors (DSPs); customized processors such as Network Processors (NPs) or Network Processing Units (NPUs), Graphics Processing Units (GPUs), or the like; Field Programmable Gate Arrays (FPGAs); and the like along with unique stored program instructions (including both software and firmware) for control thereof to implement, in conjunction with certain non-processor circuits, some, most, or all of the functions of the methods and/or systems described herein. Alternatively, some or all functions may be implemented by a state machine that has no stored program instructions, or in one or more Application-Specific Integrated Circuits (ASICs), in which each function or some combinations of certain of the functions are implemented as custom logic or circuitry. Of course, a combination of the aforementioned approaches may be used. For some of the embodiments described herein, a corresponding device in hardware and optionally with software, firmware, and a combination thereof can be referred to as “circuitry configured or adapted to,” “logic configured or adapted to,” etc. perform a set of operations, steps, methods, processes, algorithms, functions, techniques, etc. on digital and/or analog signals as described herein for the various embodiments.

(57) Moreover, some embodiments may include a non-transitory computer-readable storage medium having computer-readable code stored thereon for programming a computer, server, appliance, device, processor, circuit, etc. each of which may include a processor to perform functions as described and claimed herein. Examples of such computer-readable storage mediums include, but are not limited to, a hard disk, an optical storage device, a magnetic storage device, a Read-Only Memory (ROM), a Programmable Read-Only Memory (PROM), an Erasable Programmable Read-Only Memory (EPROM), an Electrically Erasable Programmable Read-Only Memory (EEPROM), Flash memory, and the like. When stored in the non-transitory computer-readable medium, software can include instructions executable by a processor or device (e.g., any

type of programmable circuitry or logic) that, in response to such execution, cause a processor or the device to perform a set of operations, steps, methods, processes, algorithms, functions, techniques, etc. as described herein for the various embodiments.

(58) Example User Device Architecture

(59) FIG. 3 is a block diagram of a user device **300**, which may be used in the cloud-based system **100** or the like. Specifically, the user device **300** can form a device used by one of the users **102**, and this may include common devices such as laptops, smartphones, tablets, netbooks, personal digital assistants, MP3 players, cell phones, e-book readers, IoT devices, servers, desktops, printers, televisions, streaming media devices, and the like. The user device **300** can be a digital device that, in terms of hardware architecture, generally includes a processor **302**, I/O interfaces **304**, a radio **306**, a data store **308**, and memory **310**. It should be appreciated by those of ordinary skill in the art that FIG. 3 depicts the user device **300** in an oversimplified manner, and a practical embodiment may include additional components and suitably configured processing logic to support known or conventional operating features that are not described in detail herein. The components (**302**, **304**, **306**, **308**, and **310**) are communicatively coupled via a local interface **312**. The local interface **312** can be, for example, but not limited to, one or more buses or other wired or wireless connections, as is known in the art. The local interface **312** can have additional elements, which are omitted for simplicity, such as controllers, buffers (caches), drivers, repeaters, and receivers, among many others, to enable communications. Further, the local interface **312** may include address, control, and/or data connections to enable appropriate communications among the aforementioned components.

(60) The processor **302** is a hardware device for executing software instructions. The processor **302** can be any custom made or commercially available processor, a CPU, an auxiliary processor among several processors associated with the user device **300**, a semiconductor-based microprocessor (in the form of a microchip or chipset), or generally any device for executing software instructions. When the user device **300** is in operation, the processor **302** is configured to execute software stored within the memory **310**, to communicate data to and from the memory **310**, and to generally control operations of the user device **300** pursuant to the software instructions. In an embodiment, the processor **302** may include a mobile-optimized processor such as optimized for power consumption and mobile applications. The I/O interfaces **304** can be used to receive user input from and/or for providing system output. User input can be provided via, for example, a keypad, a touch screen, a scroll ball, a scroll bar, buttons, a barcode scanner, and the like. System output can be provided via a display device such as a Liquid Crystal Display (LCD), touch screen, and the like.

(61) The radio **306** enables wireless communication to an external access device or network. Any number of suitable wireless data communication protocols, techniques, or methodologies can be supported by the radio **306**, including any protocols for wireless communication. The data store **308** may be used to store data. The data store **308** may include any of volatile memory elements (e.g., random access memory (RAM, such as DRAM, SRAM, SDRAM, and the like)), nonvolatile memory elements (e.g., ROM, hard drive, tape, CDROM, and the like), and combinations thereof. Moreover, the data store **308** may incorporate electronic, magnetic, optical, and/or other types of storage media.

(62) The memory **310** may include any of volatile memory elements (e.g., random access memory (RAM, such as DRAM, SRAM, SDRAM, etc.)), nonvolatile memory elements (e.g., ROM, hard drive, etc.), and combinations thereof. Moreover, the memory **310** may incorporate electronic, magnetic, optical, and/or other types of storage media. Note that the memory **310** may have a distributed architecture, where various components are situated remotely from one another, but can be accessed by the processor **302**. The software in memory **310** can include one or more software programs, each of which includes an ordered listing of executable instructions for implementing logical functions. In the example of FIG. 3, the software in the memory **310** includes a suitable

operating system **314** and programs **316**. The operating system **314** essentially controls the execution of other computer programs and provides scheduling, input-output control, file and data management, memory management, and communication control and related services. The programs **316** may include various applications, add-ons, etc. configured to provide end user functionality with the user device **300**. For example, example programs **316** may include, but not limited to, a web browser, social networking applications, streaming media applications, games, mapping and location applications, electronic mail applications, financial applications, and the like. In a typical example, the end-user typically uses one or more of the programs **316** along with a network such as the cloud-based system **100**.

(63) CASB System

(64) FIG. **4** is a network diagram of a CASB system **400**. The CASB system **400** can be located between the cloud-based system **100** and one or more SaaS providers **402**. As described herein, the SaaS providers **402** can be referred to as cloud providers, cloud service providers, service providers, etc. Examples of the providers **402** include, without limitation, Office 365, Dropbox, Box, Google Drive, Salesforce, etc. That is the providers **402** can provide cloud services for enterprises related to file sharing, document management, email, collaboration, scheduling, timekeeping, financial, etc. The key point is the enterprise IT is moving from local applications hosted and maintained within the enterprise network to cloud-based solutions where the data is located off-site, in the providers **402**.

(65) The CASB system **400** can be implemented in a cloud-based system, such as using the architecture of the cloud-based system **100**. The CASB system **400** can be implemented in a private cloud, a public cloud, or a hybrid cloud. Alternatively, the CASB system **400** can be one or more servers **200** that can be located on-premises with an enterprise, off-premises, etc. Even further, the CASB system **400** can be collocated with the SaaS providers **402**. That is, various architecture implementations are contemplated. Further, the CASB system **400** contemplated both operations with the cloud-based system **100**, operating as a distributed security system, as well as independent operation (i.e., with the components of the cloud-based system **100** omitted in FIG. **4**, and with the functionality incorporated in the CASB system **400** itself).

(66) The objective of the CASB system **400** is to provide enterprise IT control over data (resources) in the SaaS providers **402**. Note, as described herein, the enterprise can be referred to as a tenant of the provider **402**. The CASB system **400** is configured to operate as a distributed file crawler for files associated with a particular tenant. The CASB system **400** can both provide a report based on the file crawling as well as implement policy actions based on policy configuration.

(67) The CASB system **400** includes one or more APIs **410**, such as a Representational state transfer (REST) API. In an embodiment, the APIs **410** connect to the cloud-based system **100**, such as one of the enforcement nodes **150**. Here, a user can interact with the CASB system **400** via a User Interface (UI) **412** through the central authority **152**. Additionally, the enforcement node **150** can connect to a log **414**, such as a data store that stores statistics and transactions, for reporting. The enforcement node **150** can also connect to a DLP engine **416** for data leakage protection through the CASB system **400**. Here, the CASB **400** can be used to identify content, files, etc. that match sensitive data in a DLP dictionary. The user can provide policy and configuration via the UI **412**.

(68) Again, the CASB system **400** can be deployed without the cloud-based system **100**. Here, the API **410** can connect directly to the UI **412**, and the log **414** and the DLP engine **416** can be incorporated directly in the CASB system **400**, or in an external system.

(69) The CASB system **400** includes an authentication provider **420** that is configured to perform authentication of the tenant with the SaaS providers **402**. The APIs **410** and the authentication provider **420** connect to a message broker **422**, which is configured to interact between the APIs **410**, the authentication provider **420**, and a plurality of workers **430**. A regulator **424** is connected to the message broker. The message broker **422** is a pipeline where job tickets are queued for

consumption by the workers **430**.

(70) In an embodiment, the authentication provider **420**, a controller for the APIs **410**, the regulator **424**, and the workers **430** are Java Spring services, and other embodiments are also contemplated. The message broker **422** can be a queuing service, such as using Apache Kafka, Microsoft EventHub, or other embodiments. The API controller is a liaison service that interfaces between the CASB system **400** and the cloud-based system **100**.

(71) With respect to the authentication provider **420**, customer information, including tokens and credentials are not stored permanently or persisted. Also, the CASB system **400** is not tied specifically to a particular SaaS provider **402**. That is, the CASB system **400** is configured to operate with multiple, different SaaS providers **402**. This is accomplished through customized APIs and configuration of the workers **430**. Each SaaS provider **402** can have a different set of APIs and functionality.

(72) The workers **430** are connected to the SaaS providers **402** and are dedicated to performing particular tasks. In a sense, the plurality of workers **430** are organized in a pool of workers, and tasks are assigned between the workers **430**. The CASB **400** can include a sandbox **440** that can be connected to the DLP engine **416**, and the DLP engine **416** can also include a REST API **445** connection to the SaaS providers **402**. Note, the sandbox **440** can be included in the CASB system **400**, or it can be an external system. The sandbox **440** is configured to execute files, open files, etc. in a safe environment to analyze whether the files are malicious or not.

(73) The worker pool is a collection of workers **430** that interact with the SaaS provider **402** and perform specific tasks. The pool of workers **430** enables the CASB system **400** to operate efficiently in a distributed nature. The workers **430** are assigned tasks from various queues, via the message broker **422** and the regulator **424**. Thus, the workers **430** can operate akin to an assembly line, and there can be hand-offs between workers **430**. For example, the workers **430** can include authentication workers to authenticate users, tenants, etc., metadata workers to analyze file or content metadata, file workers to scan/analyze files, action workers to perform various policy-related actions, and the like.

(74) The workers **430** can logically be viewed as contract workers in a factory, on an assembly line, etc. The workers **430** are provided specific instructions in a job ticket. The job ticket has information on what job to be performed, where to get the inputs, and where to send the outputs. Every worker **430** also knows what to do when something goes wrong.

(75) The regulator **424** is like the SCADA (Supervisory Control and Data Acquisition) in a control system architecture. The regulator **424** monitors the performance of all the workers **430** and controls the overall system for optimum throughput.

(76) Job Ticket Example

(77) Again, the message broker **422** assigns jobs to the workers **430**. Here is an example of a job ticket for an example job:

(78) TABLE-US-00001 { TenantID : 123456 TransactionID : 111111 JobType : GetTenantUsers
Run ID : 1 SaaSProvider : Google Drive }

Design Constraints

(79) Again, each different SaaS provider **402** can have a different set of APIs and functionality. The CASB system **400** is configured to interface with a plurality of different SaaS providers **402**. The log **414** can be configured to store changes/events for an entire organization, including on a per user basis.

(80) The APIs between the CASB **400** and the SaaS providers **402** may be limited, e.g., throttled by the SaaS providers **402**. As such, there is an initial baseline crawl (i.e., a first-run) where the CASB system **400** has to crawl and scan all files in the SaaS provider **402**. This initial baseline crawl is performed efficiently and is synchronized with the DLP engine **416**. After the baseline crawl, subsequent crawls are performed incrementally, namely through files that changed since the previous crawl. For example, the first run can be referred to as run one, and each incremental crawl

is run X, which only scans and crawls files that have changed since run X-1. In an embodiment, the period of incremental calls is once a day. Of course, other periods are also contemplated.

(81) File Crawl

(82) The SaaS providers **402** generally provide two ways to crawl through the files for a tenant, namely crawling based on organization-wide file activity or a change log and crawling based on a pseudo-breadth-first traversal. The file activity or a change log enables crawling based on file changes. The pseudo-breadth-first traversal is crawling based on snapshots.

(83) FIG. 5 is a functional block diagram of filing crawling of the SaaS provider **402** with the CASB system **400**. Specifically, FIG. 5 illustrates functionality associated with file crawling in the SaaS provider **402** by the CASB **400**. The CASB **400** includes a controller **450**, such as the message broker **422** and the regulator **424**. The controller **450** can communicate with the cloud-based system **100** and the authentication provider **420**. The authentication provider **420** can communicate with the SaaS providers **402**. The CASB **400** can also include a CASB client **460** that includes a worker for DLP **462** and the log **414**. In the example of FIG. 5, there are edge workers **430a** that interface between the authentication provider **420**, the SaaS provider **402**, the controller **450**, and the CASB client **460**. The objective of the edge workers **430a** is to perform file crawling of the SaaS providers **402**. In an embodiment, the SaaS providers **402** can be file storage providers, such as, for example, Office 365 (SharePoint), Box, DropBox, etc.

(84) For illustration, an example operation is described in FIG. 5. There is a tenant event (S1) from the controller **450** to the edge worker **430a**. The next run notification (S2) is provided from the edge worker **430a** after all files are crawled in the run. The edge worker **430a** notes a new event (S3) with file meta-data, the edge worker **430a** fetches file details and provides a file for scanning (S4) which is sent to DLP **562** for scanning and analysis. A policy action (S5) can be the result of the DLP **562** and provided to the edge worker **430a**. The edge worker **430a** can implement the policy action in the SaaS provider **402** and provide the result (S6) for the log **414**. For example, a policy action can be to delete a file, quarantine a file, flag a file, etc.

(85) Crawling Based on a Change Log

(86) FIG. 6 is a flowchart of a file crawling process **500** based on a change log. The file crawling process **500** contemplates implementation by the CASB system **400** to crawl the SaaS provider **402**. The file crawling process **500** includes, for a first run (step **501**), fetching admin logs for file-related activities for a tenant in batches (step **502**), processing the batch for unique file entries in the batch (step **503**), pushing the file info into a queue (Q) (step **504**), repeating steps **503**, **504** until the entire log is crawled (step **505**), and storing the log's stream-position for a next Run (step **506**).

(87) For a run X (step **501**) where X is an integer greater than 1, the file crawling process **500** includes, fetching admin logs from the last stream-position for a tenant (step **507**), processing the batch for unique file entries in the batch (step **508**), pushing file info to a queue (Q) (step **509**), repeating through above steps **508**, **509** for all tenants until the entire log is crawled (step **510**), and storing the log's stream-position for a next Run (step **511**).

(88) Crawling Based on the Breadth-First Traversal

(89) FIG. 7 is a flowchart of a file crawling process **550** based on breadth-first traversal. The file crawling process **550** contemplates implementation by the CASB system **400** to crawl the SaaS provider **402**. For example, some SaaS providers **402** may not maintain a change log for a tenant, but instead, provide a snapshot of a user's filesystem and then a change log for every user. The file crawling process **550** includes, for a first run (step **551**), crawling through the entire list of entities (Users/Groups/SharePoint Sites) for a tenant (step **552**), and storing the list of entities for each tenant against and update Run # (step **553**). For each entity, the file crawling process **550** includes, crawling through the File System and capturing the list of files (step **554**), storing the last delta link for every entity (step **555**), and pushing the files in a queue (Q) (step **556**). The file crawling process **550** includes, after the last user, the last file pushed in the queue (Q), updating tenant info about Run # completion (step **557**), and repeating through the above steps for all tenants (step **558**).

(90) For run X (step 551) where X is an integer greater than 1, the file crawling process 550 includes fetching a list of entities for the tenant from a store (step 559), for each entity, crawling through the File System and capture the list of files (step 560), storing the last delta link for every entity (step 561), pushing the files in a queue (Q) (step 562), after the last user last file pushed in the queue (Q), updating tenant info about Run # completion (step 563), and repeating through above steps for all tenants (step 564).

(91) Flow Diagram

(92) FIG. 8 is a flow diagram of example operations between the CASB client 460, the controller 450, the message broker 422, a worker 430, and the SaaS provider 402. A new configuration or reconfiguration is provided, via the CASB client 460, the cloud-based system 100, etc. (step 601), and organization (tenant) information and credentials are provided to the controller 450 (step 602). The controller 450 gets events and pushes them in a queue (Q). The message broker 422 is configured to dequeue (D Q) the events and assign it to the worker 430 (step 604). The worker 430 is configured to interact with the SaaS provider 402 to get admin events (step 605), which are provided as JavaScript Object Notation (JSON) events (step 606.1). The process is continued until the queue is emptied, the last event in the queue (step 606.2, 606.3).

(93) The worker 430 can add new events in the queue, and the broker 422 can dequeue the new events when assigning back to a worker 430 (step 607). The worker 430 gets file info (step 608) and receives JSON file info from the SaaS provider 402. The worker 430 can scan each file in the queue (step 609), provide results to the controller 450, which dequeues the scanned file (step 610).

(94) The controller 450 can provide results of the scan to the CASB client 460, which returns information (step 611). The controller 450 can create a scan file (step 611.1) and receive a post-action (step 611.2) from the CASB client 460. For example, the CASB client 460 may perform DLP, and the action can be allow, delete, quarantine, etc. The controller 450 can implement the policy action in the queue (step 612), the brokers 422 can dequeue the policy action (step 613) and assign the action to the worker 430 which posts the action in the SaaS provider (step 614). The worker 430 can provide the action result in a queue (step 615), the broker 422 can dequeue the action results (step 616) and post the action result in the log (step 617).

(95) Webhook Integration

(96) A webhook in web development is a method of augmenting or altering the behavior of an application with custom callbacks. These callbacks may be maintained, modified, and managed by third-party users and developers who may not necessarily be affiliated with the originating website or application. Webhooks are user-defined Hypertext Transfer Protocol (HTTP) callbacks that can be triggered by some event, such as in a SaaS application detecting modification of content. When that event occurs, the SaaS application makes an HTTP request to the URL configured for the webhook. Users can configure them to cause events on one site to invoke behavior on another. SaaS applications from the SaaS providers 402 are configured to support webhooks. Webhooks operation can be compared to APIs. In APIs, one pulls data from a provider. In Webhooks, the SaaS provider 402 pushes data out, e.g., the present disclosure utilizes webhooks for identifying real-time modifications of content in SaaS applications.

(97) The present disclosure is described with reference to the CASB system 400, including all of the various architecture implementations described herein. Again, the CASB system 400 can be provided using the cloud-based system 100, in another cloud (e.g., private cloud, public cloud, hybrid cloud, etc.), as one or more servers 200 (e.g., an appliance located on-premises with an enterprise, off-premises, etc. as well as collocated with the SaaS providers 402). Note, the CASB system 400 as described herein is configured to provide CASB functionality, which may also be referred to as a CASB service. Again, the CASB service functions between the SaaS service users and the SaaS applications from the SaaS providers 402. Among the many services, CASB provides the most critical ones include identify and protect malware and policy enforcement (DLP).

(98) The present disclosure contemplates an organization (i.e., a tenant, a corporation, an

enterprise, etc.) having multiple SaaS applications and multiple users spread across multiple countries, states, cities, etc. (generally geography) providing protection as soon as possible with the CASB system **400**. That is, the CASB system **400** is configured to scan data, prevent data loss, etc. at near real-time for critical protection. To that end, the CASB system **400** is configured to detect changes in data as soon as a user **102** is active and modifies the data. Further, the CASB system **400** is configured to scan data closest to the source, using geolocation, and scanning in compliance with local law and regulations.

(99) CASB System with Webhooks Integration

(100) FIG. **9** is a flow diagram of an architecture of a CASB-webhooks system **700**, and FIG. **10** is a flowchart of a CASB-webhooks integration process **702**, which may be implemented through the CASB-webhooks system **700**, or in other approaches. The CASB-webhooks system **700** and the CASB-webhooks integration process **702** operate with the constraint that there is no replication of user data. The CASB-webhooks system **700** includes various components including a webhook controller **704**, a gateway controller **706**, an initiate subscription component **708**, a data stream processing system **710** such as Apache Kafka, a webhook listener component **712**, a subscription helper component **714**, a service publish component **716**, and a process notification component **718**.

(101) Note, these various components **704-718** can be physical or logical components that are part of the CASB-webhooks system **700** and perform the CASB-webhooks integration process **702**. The components **704-718** can be executed on one or more servers, including physical servers, virtual machines (VM) executed on physical hardware, etc. The components **704-718** can be integrated into one another, and FIG. **9** is illustrated to show functional operation. Those skilled in the art will recognize various approaches are contemplated. Also, in an embodiment, the components **704-718** can be separated from the CASB system **400**. In another embodiment, some or all of the components **704-718** can be included in the CASB system **400**. In a further embodiment, some or all of the components **704-718** can be included with the SaaS providers **402**. Of course, a combination of any of these approaches is also contemplated. Also, any of the components **704-718** may include one of the worker **430** in the CASB system **400**, as well as a worker **430** outside of the CASB system **400**.

(102) The CASB-webhooks integration process **702** includes registration (step **702-1**), renewal (step **702-2**), notifications (step **702-3**), and triggering (step **702-4**). Various, the notifications step **702-3** and the triggering step **702-4** can be implemented by the components **704, 706, 716, 710, 712, 718, 710** in communication with the SaaS provider **402** and the CASB system **400** (bottom half of FIG. **9**), and the components **708, 706, 710, 712, 714** in communication with the CASB system **400** (top half of FIG. **9**) can be used to manage the subscriptions, i.e., the registration step **702-1** and the renewal step **702-2**.

(103) Generally, the registration step **702-1** involves identifying the users **102** of a tenant, and this can be a process where IT admin performs configuration through the CASB system **400**. The registration step **702-1** can further include specifying monitoring functions per user, per group, per tenant, etc.

(104) FIG. **11** is a flow diagram of subscription and renewal process **750** for the registration step **702-1** and the renewal step **702-2**. The subscription and renewal process **750** includes a tenant (or a provider) providing a configuration (step **750-1**). This can be through the central authority **152**, the CASB client **460**, etc. The configuration can include a tenant identifier, user identifiers, actions for monitoring, content for monitoring, etc. The configuration is provided to a subscription controller (step **750-2**). The subscription controller can publish the configuration to a subscription queue (step **750-3**), which connects to a consumer subscription listener (step **750-4**) that connects to a subscription service client (step **750-5**). The subscription service client interfaces the SaaS provider **402** via batch API calls and receives responses for managing the subscriptions and configuration. This data can be stored in a database **754**.

(105) Generally, the renewal step **702-2** involves periodically renewing the subscriptions based on the timing of the SaaS provider **402**. The renewal step **702-2** operates based on a timer for each different SaaS provider **402** from a scheduler (step **750-6**). The scheduler publishes to the subscription queue (step **750-7**), which connects to the consumer subscription listener (step **750-8**). The consumer subscription listener, which can be one of the workers **430**, is configured to get subscriptions that need to be renewed and to make batch API patch requests for renewal to the SaaS provider **402** (step **750-9**). For subscription and renewal, the interaction between the SaaS provider **402** can be via controller **706**.

(106) In various embodiments, the scheduler can be further configured to control historic and ongoing scans of content of the plurality of users. Again, the scan includes identifying malware in content in the SaaS application and performing Data Loss Prevention (DLP) in the content in the SaaS application, maintaining records associated with a plurality of incidents for the malware and the DLP.

(107) Referring back to FIG. **9**, generally, the components **708**, **706**, **710**, **712**, **714** communicate with the CASB system **400** for managing the subscriptions. The notifications step **702-3** includes receiving multiple notifications via webhooks that certain users and modifying certain content. The notifications can identify the user, the content, the event type (save, delete, add, etc.). The webhook notification is provided from the SaaS provider **402** based on the subscriptions to the webhook controller **704**, which connects to the gateway controller **706**, which publishes the notifications (service publisher **716**) to the data stream processing system **710**. The webhook listener **712** detects notifications from the data stream processing system **710** and causes a process notification **718**, which can also be provided to the data stream processing system **710**, for action by the CASB system **400**.

(108) The triggering step **702-4** generally includes taking the notifications from the data stream processing system **710** that were caused by webhooks and acting upon them in the CASB system **400**. This action can include any of the monitoring and scanning functions described herein. Of note, the notifications step **702-3**, and the triggering step **702-4** can be used by the CASB system **400** for detection of which users **102** to process, to identify the event type and process with delay or process instantly, etc. Further, the notifications step **702-3**, and the triggering step **702-4** can be used by the CASB system **400** to identify which queue to push into.

(109) Geolocation

(110) FIG. **12** is a map illustrating an example system including two CASB systems **400A**, **400B** geographically distributed, and two SaaS applications **800A**, **800B** also geographically distributed. In this example, the CASB system **400A**, **400B** are single nodes in a composite CASB system **400**. Scanning data via one of the CASB systems **400A**, **400B** includes downloading the data and performing policy actions or malware detection. The SaaS application **800** users in an organization can be located anywhere, and they download and upload data while using the SaaS applications **800**.

(111) Similarly, the SaaS applications **800** can distribute and store data across the globe. Thus, strategically downloading data is critical for fast actions and remediation. To achieve this, the present disclosure includes the CASB system **400** scanning data closest to the source, which is most of the time near the location of the user. For implementation, the CASB system **400** detects the geolocation of the users **102** via their user devices **300** and routes the scan requests to the closest CASB scanners, namely the CASB systems **400A**, **400B**. Geolocation of the users **102** can be fetched periodically.

(112) Historical and Live Scanning Process

(113) FIG. **13** is a flowchart of a historical and live scanning process **850** for CASB functionality. Tenants (customers) require scanning of historical data as well as live data that is being modified by users. Historical data is scanned, with the CASB system **400**, by crawling the SaaS application using APIs. As described herein, the entities to be scanned are pushed into queues to be scanned.

Along with historical data, the present disclosure can process live data modifications using webhooks. Notifications on changes are received, as described herein, and pushed into the queues. The CASB system **400** can then prioritize the live modified data and perform the CASB functions. Specifically, historical scans are important, but they can tolerate latency. Live data modification requires scanning at near real-time speeds to detect problems and to limit the impact on the user **102** (i.e., the user **102** does not want to wait for each file, etc.).

(114) The historical and live scanning process **850** can be implemented as a method, in a server **200**, and as computer-readable code stored in a non-transitory computer-readable storage medium. The historical and live scanning process **850** includes causing a scan by the CASB system of a plurality of users associated with a tenant in a Software-as-a-Service (SaaS) application where the scan includes any of identifying malware in content in the SaaS application and identifying confidential data in the content in the SaaS application (step **851**); during the scan which is covering historical data in the SaaS application, receiving notifications of the content being actively modified by any of the plurality of users (step **852**); and including the content being actively modified in the scan with the historical data (step **853**).

(115) The historical and live scanning process **850** can further include maintaining geolocation of the any of the plurality of users (step **854**); and causing the content being actively modified in the scan to be processed by the CASB system based on the geolocation (step **855**). The historical and live scanning process **850** can further include prioritizing the content being actively modified in the scan higher than the scan of the historical data. The historical data can be scanned via Application Programming Interfaces (APIs) associated with the SaaS application, and the notifications of the content being actively modified are via webhooks from the SaaS application.

(116) The historical and live scanning process **850** can further include causing an action in the SaaS application based on the scan and based on policy and the content. The action can include any of allowing a file, deleting a file, quarantining a file, and providing a notification. The historical and live scanning process **850** can further include causing the execution of a file of the content in a sandbox for the identifying malware. The historical and live scanning process **850** can further include causing queueing of the content being actively modified and the historical data.

Additionally, in response to quarantining a file, the present systems and methods can tombstone the file which is quarantined for better user experience.

(117) CASB In-Memory Data Store

(118) As described herein, the CASB system **400**, such as implemented via the cloud-based system **100**, is configured to perform scans of data located in the SaaS providers **402**. Again, the data can include files, email, etc., and the data can be accessed by the CASB system **400** via APIs. The scan is for security and/or DLP policy, configured by a tenant. As described herein, if there is a DLP violation and/or if a file/email contains malware, this is referred to as a CASB incident or simply an incident. The present disclosure provides logging and reporting for incidents (as opposed to entire scan results). For CASB incidents, a file in the SaaS provider **402** can be modified and rescanned again and again multiple times. The logging and reporting requirement for incidents are different from the regular weblog and firewall log where the previous transactions are immutable, while incidents change. The objective here is to provide tenant IT with the current snapshot of the incidents for analysis thereof. The CASB in-memory data store describes a highly scalable and efficient incident reporting approach for the latest snapshot for a scan (for each file and email). There is a single log line recorded for each file/email that was updated when that particular file/email was last scanned.

(119) FIG. **14** is a block diagram of a CASB in-memory data store system **900**. The system **900** includes computing clusters **902** with various data nodes **904**. The data nodes **904** are where the actual data store resides. A company (tenant) can reside in one or more clusters **902** at the same time. Also, a company can historically reside in one cluster **902** and migrate to another cluster **902** as part of load balancing for the clusters **902**.

(120) The CASB in-memory data store system **900** includes data forwarder/router nodes **906** and query forwarder/merger nodes **908**. The data forwarder/router nodes **906** are configured to route logs, obtained during a scan with the CASB system **400** and based on an incident, to appropriate cluster **902**. The data forwarder/router nodes **906** perform the routing function for incidents. Note, the data forwarder/router nodes **906** have intelligence, i.e., processing capability to receive a log and to determine which cluster **902** to forward it to. Note, the CASB system **400** can be multi-tenant, and the data forwarder/router nodes **906** know the topology for each tenant, and this is used to forward a particular log, such as based on metadata identifying a tenant, to the cluster or clusters for that particular tenant. Also, if a company (tenant) is spread across multiple clusters **902**, the data forwarder/router nodes **906** can be configured to provide a log to all of the multiple clusters **902**, such as via metadata. For example, in FIG. **14**, assume there are two clusters **902** A and B, if a tenant is on both, the data forwarder/router nodes **906** attempt to distribute the log (record) equally among each cluster **902**. In an embodiment, the receives the log for each incident from a worker **430** in the CASB system **400**.

(121) The query forwarder/merger nodes **908** are configured to receive queries from an external source and performs such queries through the data nodes **904**. All requests (queries) land on query forwarder/merger nodes **908** and a node, based on a metadata of the tenant to cluster mapping, performs query planning, sends requests to clusters **902**, and, if the query has been sent to multiple clusters **902**, it will merge the results and send responses back to the client (requestor). Similar to the data forwarder/router nodes **906**, the query forwarder/merger nodes **908** have the state to route the queries to a particular cluster **902**, including multiple clusters **902** where a tenant is spread. Once the query forwarder/merger node **908** receives a response, it merges the results if there where multiple clusters **902**, and sends the result back to the client.

(122) Data Store Design and Data Structure

(123) The data store in the CASB in-memory data store system **900** supports various operations including an insert operation, a delete operation, an update operation, and a recovery mechanism to rebuild the data store in case of a service restart. FIG. **15** is a flowchart of a record processing process **950** implemented in the data store. The process **950** describes how update, delete, and insert operations are performed.

(124) A new record is received (step **951**). In case of any new record (step **951**), the process **950** includes checking if entry exists in a filehash in a data store **960** (step **952**). If so (step **952**), it means the process **950** has seen that file earlier, and the old record is deleted from the filehash and deleted from the data store in a deferred delete, described herein (step **953**). If the file does not exist in the filehash in the data store **960** or after the delete step **953**, the record is inserted in the filehash in the data store **960** and in the data store (step **954**). For a recovery mechanism, the process **950** can include periodically flushing the data store logs to a file (step **955**). This can be every so often, based on a trigger event, etc.

(125) The data store is mutable in-memory for highly-optimized updates and reports, for the most common dimensions of filtering, aggregation, and pagination. In the CASB system **400**, the most common dimensions are application name (appname or application ID) and tenant name (i.e., the user **203**). As such, all reports are around these dimensions.

(126) FIG. **16** is a diagram of an example implementation of the filehash in the data store **960**. The filehash in the data store **960** is stored in the data store and can be a multilevel hash that includes a hash of a customer ID, application ID, tenantname ID, etc., where data is hashed by company, application name (appname or appid), and tenant name (tenantname or tid). At the leaf level, a list of actual records is stored. This data store will always contain the latest incident information for any file and for a particular combination of (appname, tenant, filename), only one entry is present, referred to as a datanode in FIG. **16**.

(127) Also, the following facts are noted and assumptions made. A file or email that is being scanned can belong to only one tenant (user), and a tenant can belong to one application. The

maximum scan duration supported by the data store is for a certain period of time. Older scan results will be cleaned up periodically by the product. Most reports come with company, appname and tenant filters. Majority of the aggregation queries are on the filenames/email. So, the data store optimizes these reports.

(128) The following provides pseudo code for the filehash in the data store **960**: 1) New rec received 2) Check if already same file already exist for cid,appid, tid combination In filehash

```
(129) TABLE-US-00002 filehash = smhashapi_keymd5_old(filenamep, filelength);  
head=&smslave_file_hash_table[(datanode->filehash)%SMSLAVE_FILEHASH_SIZE] ;  
_TAILQ_FOREACH(entry, head, filehash_entry) { If (cid, appid and tid matches) /* delete  
old record if it is not referenced both from this table as well as from inmemory /* data store  
} 3) #define PREPARE_KEYREC(_cid, _appcat, _appid, _tid) keys[keycnt++]._u_int=_cid;  
keys[keycnt++]._enum=_appid; keys[keycnt++]._u_int=_tid;
```

This will be used to find or insert record at which leaf node that is to be put in the in-memory data store 4) Find leaf node for zks_prefix_tree_schema_get(ctx->ptp_dstore, keys) for keys prepare in step (3)

This will give us the head of list where actual complete record is stored.

```
5)_TAILQ_INSERT_HEAD(head, datanode, comp_list_entry
```

This will make sure records are stored in descending order of time.

Data Duplication Prevention

(130) In the CASB in-memory data store system **900**, data is distributed across the nodes **904** in clusters **902** based on some hash functions (e.g., key%noOfnodes) and records based on some key that can also be updated/deleted in mutable systems. But when any node **904** is added or deleted, the hash function can give different nodes for a key which was earlier sent to node A and later the hash function sends to another node B leading to multiple copies of the same mutable record being on different nodes **904**, leading to inconsistency in data. This scenario can occur when a company moves from one cluster to another cluster as part of the load balancing.

(131) FIG. 17 is a diagram of data duplication prevention in the data store. To make sure that the only one record of a key lies across all the nodes **904** in a cluster **902**, whenever any node addition happens, a broadcast is sent to delete that key to all of the nodes **904** in the cluster **902** except the current node which the hash function distributes to. In that way, all nodes **904** which have an old record of same key will delete their record and only one node will have the record.

(132) Suppose in the CASB in-memory data store system **900**, there is data with a key as FILEID distributed uniquely among 3 nodes (Node 1, 2, 3). FileId 4 belongs to Node 2. Now, Node 4 is added and due to this, the data forwarder forward it to Node 4 leading to multiple copies of FILEID 4. During this time, the data forwarder/router node **906** detects that node addition has been done in last X days and there can be chances that this data might be lying on other nodes, the data forwarder/router node **906** sends a DELETE record of that key (in this case FILEID 4) to all other nodes which might be having this older record and each node manager handles the delete record in Deferred delete way which is covered next.

(133) Deferred Delete on Data Manager

(134) Data managers cannot delete a record blindly without checking a reference for that record otherwise it will show inconsistent record when query/report runs. A deferred delete includes, suppose a query is running and referencing the data store, and a delete record is received where a query is also running. The data manager checks if the record is referred by a query manager, and if yes, puts that record into a deferred list and marks it a soft delete. Whenever the reference of that record becomes zero, the data manager deletes that record.

(135) For insertion, it is checked from another file-based hash whether an entry is already there in the data store, and, if yes, delete that record using the deferred delete operation and insert the new one. Insertion is fast as the process just needs to find a particular tenant hash bucket and insert at the head.

(136) Benefits of Maintaining a Unique Key Across Nodes

(137) In distributed query engines, most of the expensive queries are related to count distinct queries on high cardinality fields with aggregates. Consider the use case for the following query:

(138) “from CASB select count(distinct filename) where companyid=1 group by appname, tenantid”

(139) The above query is an expensive query as it needs to maintain the list of all files in memory to resolve duplicates and count the filenames after all data has been read. Due to the nature of the data store (multi-level hash), this complex count distinct query on unique keys can be converted into a normal count query leading to faster response time. Thus, the query is converted to

(140) “from casb select count(1) as unique files where cid=1 group by appname, tid”

(141) Instead of maintaining the entries for each file, it just needs to read and count the rows as duplicates are removed at the time of record update.

(142) Reporting Design and optimizations

(143) A query request can come to query forwarder/merger node **908** for following kind of companies: a) Company which is not distributed across multiple clusters **902**. In this case, requests need to be processed from an in-memory data store lying on a single data node and the node **908** acts as a forwarder for these kinds of queries. b) Company which is distributed among multiple clusters **902**. In this case, requests need to be processed from many in-memory data store lying on different clusters **902** and the node **908** acts as query planner and merger for these kinds of queries.

(144) For all these requests, optimized filtering and aggregation operations are needed to have faster response time and reduce the resource utilization.

(145) Filtering

(146) The most common reports require filters on appname and tenant name. The data store is maintained such that if any such filter comes, it is possible to apply filtering at the hash level instead of reading all records. This will be applied at the hash level and once level filtering is passed, data can be directly passed to the consumer without applying the same filter again and again on each record.

(147) Advantageously, a chain of records is filtered by applying a filter at level of in memory data store **960** skipping filter check for each record. This technique of using in-memory data store **960** for filtering can be done for other use cases.

(148) Inbuilt Cache for Popular Reports (Group by Appname, Tenantname)

(149) The data is stored in a multi-level hash (FIG. **16**) that includes company, application and tenant nodes. These nodes are the parent to underlying child nodes that contain the incident scan results. These nodes can work as a cache to store the summary of the underlying logs. For example, each application node can store the following information: i) Count of malware incidents in the underlying nodes, e.g., from CASB select count(transactions) where malware >0 where appname=BOX group by tenant; ii) Count of DLP violation, e.g., from CASB select count(transactions) where DLP>0 group by appname, and iii) number of files in the tenant.

(150) The count on the parent nodes are updated when a new incident is added/modified in the data store **960**. The data update operation will be negligible because each file belongs to only one <company, appname, tenant> combination.

(151) Since the data store **960** supports the scan reports for the last X days (e.g., 90 days), the parent nodes can store the counts for each separate day by maintaining X buckets. All the reports starting with day boundary time can be served from the parent node cache.

(152) Pagination for Aggregate Reports

(153) In a distributed data system, running aggregation based queries on high cardinality fields can be very expensive. For example, from CASB select count(distinct filename) as trans group by ownername limit 25 offset 0. Here, the filename can have high cardinality in the range of millions, the ownername can have medium cardinality in range of 100's of thousand, and, in this query, a user wants to see 25 aggregated records.

(154) The query forwarder/merger node **908** can send this query to all the individual data nodes **904**, which can each have a query manager, and the query manager processes the query as follows: 1) start reading data and create hash based on ownername column; 2) for each hash bucket (unique ownername values), maintain unique distinct filename values; and 3) once data read is finished, take 25 bucket and count unique A values and send results.

(155) In this approach, even if the query forwarder/merger node **908** asks for 25 aggregate records, the query processor is processing all 100 k unique B entries and maintaining a hash for all the buckets leading to query slowness. Also, the hash lookup for all the values of aggregating column which will be processor intensive and maintaining hash for all the values of aggregating column will take too much memory.

(156) To solve this issue, the query forwarder/merger node **908** forwards the request to each query processor on each node **904** and asks for 25 records from each query processor. The query processor maintains a hash for 25 entries instead of maintaining a hash for all the 100K unique entries

(157) The query processor maintains metadata of the aggregated column which it has seen previously. The metadata is used to skip previous run response entries which the query processor has already sent. When this request is received the first time with offset 0, the metadata will be empty. The following steps for the first run are as follows: a) start reading data and create hash based on ownername column. b) once a hash contains 25 unique entries (limit count), the query manager creates a filter with these values and starts dropping entries for the rest of (100 k-25) entries instead of maintaining hash and doing lookup. This will help in reducing memory footprint and query faster. c) once the query is finished on the query node, each query node will send its result to the query forwarder node **908** and each query processor on that node maintains metadata of these unique entries for the next pagination request if it comes.

(158) The steps for the second run where a client requests the next 25 records: a) the query processor starts reading data and puts the entry into the hashtable only if it does not belong to metadata entries (previous run response unique entries); b) now when a hashtable contains unique 25 entries, the query manager creates a filter with these values and starts dropping entries for the rest of entries instead of maintaining hash and doing lookup; c) once the query is finished on the query node, each query node will send its result to the query forwarder and each query processor on that node updates metadata of these unique entries for the next pagination request if it comes; and d) now metadata will have 50 entries (25 for 1st run and 25 for 2nd run).

(159) The same can continue for the next pagination request.

(160) Advantageously, this reduces the memory footprint for the query engine as there is no need of doing aggregation for all the values of a dimension, and the response time will be faster.

(161) For queries spanning multiple in-memory data store nodes **904**, the query forwarder/merger node **908** will merge the results and return the result to the client. For a single in-memory data store, the query forwarder/merger node **908** will forward the result to the client.

(162) Stateless Query Execution Engine with Cached Support

(163) For all expensive queries, the results can be cached after the query execution. The cache could be used by the subsequent queries in the following cases: i) a similar query is received from another user and the cached results can be served directly to the client. ii) The cached data structure can be further processed to produce the report. In this case, the query will be faster as it will be working on a subset of data rather than a complete data store.

(164) For example, suppose the first query is “from CASB select count(distinct filenames) group by appname, tid”

(165) In this, once the query is finished, the results are sent and saved in cache with an Abstract Syntax Tree (AST) acting as metadata for query matching.

(166) For example, suppose a second query is “from CASB select count(distinct filenames) where appname=BOX group by tid”

(167) Once the query is parsed, the query engine will try to check based on incoming Query AST (Abstract Syntax Tree) by matching it previous query AST's and see if it matches. In this example, group by clause AST match, cache contains appname and tid and the query contains tid only so this query can be run as subset of original query.

(168) If the select clause also matches, where the clause of the 2nd query is also a subset of original query as original query contains data for everything and also appname field in where clause is present in group by AST of cache query (appname). In this case, the query will be faster because it is not reading all the data, it is just reading the aggregate results and from there deriving response for this query result.

(169) For example, suppose a third query is

(170) "from CASB select count(distinct filenames) where ownername like ("abhishek") group by tid"

(171) Here, the group by AST is subset of group by AST of query cache (query 1), the select clause AST is subset of select AST of query cache (query 1), and the where clause AST does not match cache AST because ownername field is not there in group by AST, select by AST of cache AST.

(172) Algorithm

(173) a) if a filter matches multilevel keys (appname or tid), it converts it into level filter; b) if it is a count distinct query on a unique key with a group by multilevel keys convert it into a normal count query.

(174) For example "from CASB select count (distinct filename) group by appname, tid," this can be converted to from CASB select count(filename) group by appname, tid. c) count based queries multi-level hash keys can be served from node count metrics instead of walking through multilevel hash leaf entries. In this, the query engine producer will be just passing the level key and metrics to consumer instead of passing each record.

Sample Queries

(175) from casb_fileapps select count(distinct filename) as tot_dlpincidents where appname=BOX and ANY(dlpdictids)!=0".fwdarw.output: tot_dlpincidents=34.

(176) from casb_fileapps select count(distinct filename) as unique_fileincidents, id2name(ownerid) as ownername group by userid ownerid limit 100".fwdarw.output
ownerid,unique_fileincidents,ownername, 65761,130,traffic_loc->other, 65786,1,ca_loc.

(177) from casb_fileapps select time as time, filename as filename, appname as appname id2name(tid) as tenantname, dlpdictids where ANY(dlpdictids)!=0 limit 2".fwdarw.output
1594175160, one.txt, BOX, ca_loc, [61,62], 1594167110, casb_test.txt, GDRIVE, g_tenant, [1,3].

(178) CASB UI

(179) The CASB UI systems and methods including onboarding, policy configuration, reporting, advanced reporting, security exceptions, etc. Each of these is described as follows with associated UI screens and actions by an IT administrator (user). The onboarding supports configuring new cloud applications for the CASB **400**, including, for example, OneDrive, SharePoint, Exchange, Box, Google Drive, Gmail, Dropbox, and the like. The policy configuration relates to configuring DLP and/or malware rules for the onboard cloud applications. The reporting and advanced reporting are utilized to provide users a uniform view of CASB operations in the onboard cloud applications (SaaS).

(180) Starting with a SaaS Data Exposure Report, the customer gets a high-level overview of the number of assets being exposed to incidents for the entire organization. Tenant-level details are shown per application with an accordion, which can be expanded or collapsed on demand. The customer could then drill down to see the report for a specific application or tenant. The application/tenant-level report is broken down into three main sections: Assets with incidents; Owners of the assets with incidents; and Collaborators of the assets with incidents.

(181) The sections could be further broken down by the incident type or collaboration scope, which allows the user to filter the numbers. The table-based report will show the detailed information of

assets/owners/collaborators. For example, for each asset, it shows the path, owner, severity, etc. The customer could easily see the scan history for this asset, in particular, to investigate what incidents it has generated.

(182) In CASB Report, organizations can get the analytics of how to manage and protect the data stored on the cloud. This will also protect and prevent the loss of sensitive data across all cloud services in an environment. This report includes a clean interface in which users will get to know how their policies are enforced.

(183) The UI includes pie, multi-line, and bar charts to complete the design to display a massive amount of data, including re-rendering based on the user actions. The system is able to effectively handle the large set of data and render the most sensible segment of data with a minimal resource utilization (browser cache, RAM).

(184) FIGS. **18-35** are screenshots of a UI associated with the CASB system **400**. The screenshots can be displayed on a display such as associated with the user device **300**, e.g., via a Browser, via a dedicated application, etc. The user device **300** can be connected to the controller **450** and/or to the cloud-based system **100**. That is, the rendering of the UI can be via the controller **450** or any of the nodes **150**, **152** in the cloud-based system **100**. The UI includes a dashboard tab, an analytics tab, a policy tab, an administration tab, an activation tab, etc.

(185) Onboarding

(186) FIGS. **18-23** are screenshots of an onboarding process for enrolling an organization through the CASB system **400** with a SaaS application. The onboarding process is under the administration tab. In an example embodiment, the CASB system **400** supports OneDrive, SharePoint, Exchange, Box, Google Drive, Gmail, and Dropbox for the SaaS applications. Of course, other SaaS applications are also contemplated. In FIG. **18**, a user (i.e., IT administrator, etc.) is presented an Add SaaS Application Tenant screen. For example, popular SaaS applications are displayed (via boxes/icons **1002**) as well as a data entry box **1002** for manually typing in the SaaS application. In FIG. **19**, a user types “Microsoft . . .” in the data entry box **1002** and is presented with three options for SaaS applications—Exchange, OneDrive, SharePoint. data entry box **1002**. In FIG. **20**, Exchange has been selected and the user supplies a tenant name **1006** and credentials **1008** for enrollment. FIG. **21** is a similar onboarding screen for Dropbox. Those skilled in the art will recognize these onboarding screens will be specific to the SaaS application. FIG. **22** is a screenshot of onboarding failure, and FIG. **23** is a screenshot of onboarding success. After the onboarding success, the user can list trusted domains and trusted users for accessing the SaaS application.

(187) Policy

(188) FIGS. **24-27** are screenshots of a policy configuration process, subsequent to the onboarding process. The policy configuration process is under the policy tab and allows configuration of DLP, malware, and scanning configuration policy. In this example, Exchange is selected as the SaaS application for illustration purposes. In FIG. **24**, the user selects the DLP configuration for Exchange including selection of one or more DLP engines and actions (report incidents only). In FIG. **25**, the user selects the malware detection for Exchange including malware protection and a quarantine location. FIG. **26** illustrates an example DLP Rule for Email application tab, and FIG. **27** illustrated an example Add Malware Rule for Email application.

(189) Reports

(190) FIGS. **28-35** are screenshots of various reports and visualizations associated with monitoring through the CASB system **400**. FIGS. **28** and **29** are a view of the analytics tab, specifically FIG. **28** can be a first top of the screen whereas FIG. **29** can be a bottom half of the screen. FIG. **28** includes a pie chart **1100** of incidents split between DLP violations and malware, a line chart **1102** that illustrates incidents over time (separate lines for all incidents, DLP incidents, and malware incidents, and a table of SaaS application illustrating onboard SaaS applications with a number of total incidents, DLP incidents, and malware incidents.

(191) FIG. **29** includes pie charts **1120**, **1122**, bar graphs **1124**, **1126**, and a table **1128** of risky

cloud applications in use. The pie chart **1120** illustrates DLP violations split by different DPI engines (e.g., customer list, source code, etc.). The pie chart **1122** illustrates malware by threat category (e.g., virus, spyware, trojans, adware, etc.). The bar graph **1124** illustrates an organization's departments arranged by the most incidents, and the bar graph **1126** illustrates the organization's users by the most incidents. The table **1128** lists risky cloud applications and their associated category (e.g., file sharing, social networking, etc.), threat index (a value indicative of risk), bandwidth statistics, and recommended remediation (caution, block, etc.).

(192) In various embodiments, granular remediation is contemplated. For example, granular remediation can include removing an external user from write, but still allow the external user to read on a single configuration. Additionally, other specific privileges can be selectively removed for specific users in addition to specific privileges being maintained for the users. These remediations can be automatically initiated by the present systems and methods described herein.

(193) FIG. **30** is the pie chart **1120** illustrating an ability to change the values. For example, the pie chart **1120** can display values by DLP engine, by DLP severity, and by DLP dictionary. FIG. **31** is the line chart **1102** illustrating an ability to change the data. Specifically, the line chart **1102** can show values over a customized time, by select or all SaaS applications, etc.

(194) FIG. **32** is a table **1130** of an insight log that lists specific CASB incidents, including event time, policy type, application, rule name, action, and tenant. FIG. **33** is a table **1140** of email messages with incidents and tiles **1142** illustrating summary values. FIG. **34** is a pop up **1150** when a user selects a specific email in the table **1140**.

(195) Asset Reports

(196) FIG. **35** is a screenshot of a SaaS summary report for an organization. This includes different SaaS applications with summary values, e.g., files with incident, files with DLP violations, files with malware, and externally shared files with DLP violations. In this example, OneDrive is expanded showing bar graphs of files per tenant. Also, the SaaS summary report includes summary values **1160** at the top with files with incident, files with DLP violations, files with malware, and externally shared files with DLP violations across all SaaS applications. Further, the SaaS summary report includes a pie chart **1162** showing violations spread across all SaaS applications.

(197) As described herein, the CASB incidents are in relation to objects, such as file, emails, email attachments, etc. FIG. **36** is a screenshot of an assets with incidents view illustrating a file tab. FIG. **37** is a screenshot of the assets with incidents view illustrating an owners tab. FIG. **38** is a screenshot of a popup view illustrating files with incidents for a specific owner. FIG. **39** is a screenshot of the assets with incidents view illustrating a collaborators tab.

(198) FIGS. **36-39** provides summaries across three object level views to answer the most common questions that IT Practitioners have to answer before during the process of discovery and assessment of Asset Exposure for DLP and compliance. For example, FIG. **36** answers these questions with tiles—files with incident, files with DLP violations, files with malware, and externally shared files with DLP violations. With respect to who owns the affected data, FIG. **37** answers with the tiles under the OWNERS tab, e.g., Owners of publicly shared files, Owners of files with unauthorized external collaborators, etc. With respect to who from the outside has access to this affected data, FIG. **39** answers with the tiles under the COLLABORATORS tab, e.g., External collaborators on assets with sensitive data in it. Answering these questions is the key value. Typically, in other solutions, practitioners have to perform a lot of offline analysis to find answers for these. The CASB Object Level Reporting described herein provides these answers right away.

(199) FIGS. **40-43** are screenshots of a security exception process for enrolling trusted organizations, users, third-parties, etc. (trusted user exceptions) through the CASB system **400** with a SaaS application. The present systems and methods allow trusted users to be identified via the CASB UI of the present disclosure. For example, FIG. **40** shows an input screen with a section for entering trusted domains and users **1164**. The present UI can present a drop down list for selecting

one or more trusted domains, shown in FIG. 41. Additionally, security exceptions can be entered for specific organizations, users, third-parties, etc. (users). The security exceptions can include any policy and rule exceptions for any user, tenant, group of users, etc. through the security exception interface shown in FIG. 42. The UI can display all or a portion of the tenants, and display which tenants include trusted domains, trusted users, etc. FIG. 43 is a screenshot of the UI of the present disclosure displaying a list of tenants including the number of trusted domains **1166** and trusted users **1168**.

(200) More specifically, the one or more specific users are generally considered external to the primary organization or enterprise. The exceptions can allow trusted users to access sensitive data hosted on SaaS applications to perform a plurality of activities such as M&A activity, onboarding of subsidiaries, external contracting agencies helping with projects, auditors or reviewers evaluating disputes, researchers contributing to product ideas, and the like. Providing exceptions for the actions described can include using single trusted domains where all users of the domain can be trusted. Additionally, trusted users can be defined as asset of identified users that are added to the exceptions. The exceptions can be created to allow third-party users to be treated like employees when their access and activities are evaluated across CASB policies.

(201) CASB UI Process

(202) FIG. 44 is a flowchart of a CASB UI process **1200**. The CASB UI process **1200** can be implemented as a method, in a server **200**, and as computer-readable code stored in a non-transitory computer-readable storage medium. The CASB UI process **1200** includes, responsive to a scan by the CASB system of a plurality of users associated with a tenant in a Software-as-a-Service (SaaS) application where the scan includes identifying malware in content in the SaaS application and performing Data Loss Prevention (DLP) in the content in the SaaS application, maintaining records associated with a plurality of incidents for the malware and the DLP (step **1202**); providing a User Interface (UI) for the tenant including an analytics view with a plurality of summary tiles including visualizations of the plurality of incidents for the malware and the DLP for the tenant (step **1204**); and providing the UI for the tenant including a table listing any of the plurality of incidents for the malware and the DLP for the tenant, including any of unique data objects, unique users internal to the tenant, and unique external entities, associated with the plurality of incidents (step **1206**).

(203) The CASB UI process **1200** can further include providing the UI for the tenant to onboard a plurality of SaaS applications including the SaaS application. The CASB UI process **1200** can further include providing the UI for the tenant to configure policies for the DLP and for the malware for the SaaS application. The CASB UI process **1200** can further include, responsive to a selection of any entry in the table, providing a popup listing details associated with the corresponding incident. The SaaS application can be one of a plurality of SaaS applications for the tenant, and wherein the visualizations include a table listing the plurality of incidents associated with the plurality of SaaS applications. The visualizations can include one or more pie charts illustrating the plurality of incidents. The visualizations can include a line chart illustrating the plurality of incidents over time.

(204) FIG. 45 is a flowchart of another embodiment of a CASB UI process **1300**. The CASB UI process **1300** can be implemented as a method, in a server **200**, and as computer-readable code stored in a non-transitory computer-readable storage medium. The CASB UI process **1300** includes providing a user interface (UI) for a tenant to input one or more malware and data loss prevention (DLP) rules, and trusted user exceptions, wherein the trusted user exceptions identify one or more specific users and rule exceptions for the specific users (step **1302**). Responsive to a scan by the CASB system of a plurality of users associated with a tenant in a software-as-a-service (SaaS) application where the scan includes identifying malware in content in the SaaS application and performing DLP in the content in the SaaS application based on the one or more malware and DLP rules and trusted user exceptions, maintaining records associated with a plurality of incidents for the malware and the DLP (step **1304**). Providing the UI for the tenant including an analytics view

with a plurality of summary tiles including visualizations of the plurality of incidents for the malware and DLP for the tenant and a table listing any of the plurality of incidents for the malware and the DLP for the tenant, including any of unique data objects, unique users internal to the tenant, and unique external entities, associated with the plurality of incidents (step **1306**).

(205) The steps can further include performing remediation of the plurality of incidents, wherein the remediation of the plurality of incidents includes granular remediation, or wherein the remediation of the plurality of incidents includes tombstoning one or more files for better user experience. A scheduler can be configured to control historic and ongoing scans of content of the plurality of users. It will be appreciated that the steps in process **1200** and process **1300** can be combined in any way to provide other embodiments of a CSAB UI process.

(206) Tombstoning one or more files can further include, responsive to scanning users files, finding files containing malware or sensitive data and automatically moving the files to quarantine or sensitive data hold. The next time the user logs in, they will not see the files, which may result in the user creating/uploading the same content, or open an IT case to report the missing data. The present disclosure avoids such issues by leaving tombstone files in place of the content that was moved to quarantine or sensitive data hold location. The tombstone can be a text file containing details such as what happened, why the file was moved, where to find the original file, acceptable data usage policy defined by the enterprise, and other content of the like.

(207) Conclusion

(208) Although the present disclosure has been illustrated and described herein with reference to preferred embodiments and specific examples thereof, it will be readily apparent to those of ordinary skill in the art that other embodiments and examples may perform similar functions and/or achieve like results. All such equivalent embodiments and examples are within the spirit and scope of the present disclosure, are contemplated thereby, and are intended to be covered by the following claims.

Claims

1. A non-transitory computer-readable storage medium having computer-readable code stored thereon for programming one or more processors associated with a Cloud Access Security Broker (CASB) system to perform steps of: providing a User Interface (UI) for a tenant to input one or more malware and Data Loss Prevention (DLP) rules, and trusted user exceptions, wherein the trusted user exceptions identify one or more specific users and rule exceptions for the specific users; responsive to a scan by the CASB system of a plurality of users associated with a tenant in a Software-as-a-Service (SaaS) application where the scan includes identifying malware in content in the SaaS application and performing DLP in the content in the SaaS application based on the one or more malware and DLP rules and trusted user exceptions, maintaining records associated with a plurality of incidents for the malware and the DLP; and providing the UI for the tenant including an analytics view with a plurality of summary tiles including visualizations of the plurality of incidents for the malware and DLP for the tenant and a table listing any of the plurality of incidents for the malware and the DLP for the tenant, including any of unique data objects, unique users internal to the tenant, and unique external entities, associated with the plurality of incidents.
2. The non-transitory computer-readable storage medium of claim 1, wherein the steps further include performing remediation of the plurality of incidents.
3. The non-transitory computer-readable storage medium of claim 2, wherein the remediation of the plurality of incidents includes granular remediation.
4. The non-transitory computer-readable storage medium of claim 2, wherein the remediation of the plurality of incidents includes tombstoning one or more files for better user experience.
5. The non-transitory computer-readable storage medium of claim 1, wherein a scheduler is configured to control historic and ongoing scans of content of the plurality of users.

6. The non-transitory computer-readable storage medium of claim 1, wherein the steps further include providing the UI for the tenant to onboard a plurality of SaaS applications including the SaaS application.
7. The non-transitory computer-readable storage medium of claim 1, wherein the steps further include responsive to a selection of any entry in the table, providing a popup listing details associated with the corresponding incident.
8. A method comprising steps of: providing a User Interface (UI) for a tenant to input one or more malware and Data Loss Prevention (DLP) rules, and trusted user exceptions, wherein the trusted user exceptions identify one or more specific users and rule exceptions for the specific users; responsive to a scan by the CASB system of a plurality of users associated with a tenant in a Software-as-a-Service (SaaS) application where the scan includes identifying malware in content in the SaaS application and performing DLP in the content in the SaaS application based on the one or more malware and DLP rules and trusted user exceptions, maintaining records associated with a plurality of incidents for the malware and the DLP; and providing the UI for the tenant including an analytics view with a plurality of summary tiles including visualizations of the plurality of incidents for the malware and DLP for the tenant and a table listing any of the plurality of incidents for the malware and the DLP for the tenant, including any of unique data objects, unique users internal to the tenant, and unique external entities, associated with the plurality of incidents.
9. The method of claim 8, wherein the steps further include performing remediation of the plurality of incidents.
10. The method of claim 9, wherein the remediation of the plurality of incidents includes granular remediation.
11. The method of claim 9, wherein the remediation of the plurality of incidents includes tombstoning one or more files for better user experience.
12. The method of claim 8, wherein a scheduler is configured to control historic and ongoing scans of content of the plurality of users.
13. The method of claim 8, wherein the steps further include providing the UI for the tenant to onboard a plurality of SaaS applications including the SaaS application.
14. The method of claim 8, wherein the steps further include responsive to a selection of any entry in the table, providing a popup listing details associated with the corresponding incident.
15. A system associated with a Cloud Access Security Broker (CASB) system, comprising: one or more processors and memory storing instructions that, when executed, cause the one or more processors to provide a User Interface (UI) for a tenant to input one or more malware and Data Loss Prevention (DLP) rules, and trusted user exceptions, wherein the trusted user exceptions identify one or more specific users and rule exceptions for the specific users; responsive to a scan by the CASB system of a plurality of users associated with a tenant in a Software-as-a-Service (SaaS) application where the scan includes identifying malware in content in the SaaS application and performing DLP in the content in the SaaS application based on the one or more malware and DLP rules and trusted user exceptions, maintain records associated with a plurality of incidents for the malware and the DLP; and provide the UI for the tenant including an analytics view with a plurality of summary tiles including visualizations of the plurality of incidents for the malware and DLP for the tenant and a table listing any of the plurality of incidents for the malware and the DLP for the tenant, including any of unique data objects, unique users internal to the tenant, and unique external entities, associated with the plurality of incidents.
16. The system of claim 15, wherein the instructions that, when executed, further cause the one or more processors to perform remediation of the plurality of incidents.
17. The system of claim 16, wherein the remediation of the plurality of incidents includes granular remediation.
18. The system of claim 16, wherein the remediation of the plurality of incidents includes tombstoning one or more files for better user experience.

19. The system of claim 15, wherein a scheduler is configured to control historic and ongoing scans of content of the plurality of users.

20. The system of claim 15, wherein the instructions that, when executed, further cause the one or more processors to responsive to a selection of any entry in the table, provide a popup listing details associated with the corresponding incident.
