

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12395843
Kind Code	B2
Date of Patent	August 19, 2025
Inventor(s)	Srinivas; Anand et al.

SASE services for private mobile network

Abstract

Some embodiments provide a method for implementing a software-defined private mobile network (SD-PMN) for an entity. At a physical location of the entity, the method deploys a first set of control plane components for the SD-PMN, the first set of control plane components including a security gateway, a user-plane function (UPF), an AMF (access and mobility management function), and an SMF (session management function). At an SD-WAN (software-defined wide area network) PoP (point of presence) belonging to a provider of the SD-PMN, the method deploys a second set of control plane components for the SD-PMN that includes a subscriber database that stores data associated with users of the SD-PMN. The method uses an SD-WAN edge router located at the physical location of the entity and a SD-WAN gateway located at the SD-WAN PoP to establish a connection from the physical location of the entity to the SD-WAN PoP.

Inventors: Srinivas; Anand (San Francisco, CA), Gao; Xiao H. (Chapel Hill, NC), Naik; Sameer (San Jose, CA)

Applicant: VMware LLC (Palo Alto, CA)

Family ID: 1000008762393

Appl. No.: 18/071543

Filed: November 29, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20240073694 A1	Feb. 29, 2024

Related U.S. Application Data

us-provisional-application US 63402057 20220829

Publication Classification

Int. Cl.: H04W12/086 (20210101); H04L9/40 (20220101); H04L12/46 (20060101); H04L12/66 (20060101); H04L41/0654 (20220101); H04L41/0806 (20220101); H04L41/0895 (20220101); H04L41/122 (20220101); H04L41/5019 (20220101); H04L45/28 (20220101); H04L45/64 (20220101); H04L45/74 (20220101); H04L45/76 (20220101); H04L45/80 (20220101); H04W8/18 (20090101); H04W12/06 (20210101); H04W12/08 (20210101); H04W12/30 (20210101); H04W16/18 (20090101); H04W24/04 (20090101); H04W36/30 (20090101); H04W40/24 (20090101); H04W76/12 (20180101); H04W84/04 (20090101); H04W88/16 (20090101)

U.S. Cl.:

CPC H04W12/086 (20210101); H04L12/4633 (20130101); H04L12/66 (20130101); H04L41/0654 (20130101); H04L41/0806 (20130101); H04L41/0895 (20220501); H04L41/122 (20220501); H04L41/5019 (20130101); H04L45/28 (20130101); H04L45/64 (20130101); H04L45/74 (20130101); H04L45/76 (20220501); H04L45/80 (20220501); H04L63/0272 (20130101); H04W8/18 (20130101); H04W12/06 (20130101); H04W12/08 (20130101); H04W12/35 (20210101); H04W16/18 (20130101); H04W24/04 (20130101); H04W36/30 (20130101); H04W40/24 (20130101); H04W40/246 (20130101); H04W76/12 (20180201); H04W84/045 (20130101); H04W88/16 (20130101)

Field of Classification Search

CPC: H04W (12/086); H04W (8/18); H04W (12/06); H04W (12/08); H04W (12/35); H04W (16/18); H04W (24/04); H04W (36/30); H04W (40/24); H04W (40/246); H04W (76/12); H04W (84/045); H04W (88/16); H04L (12/4633); H04L (12/66); H04L (41/0654); H04L (41/0806); H04L (41/0895); H04L (41/122); H04L (41/5019); H04L (45/28); H04L (45/64); H04L (45/74); H04L (45/76); H04L (45/80); H04L (63/0272); H04L (41/5009); H04L (45/22)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
9450817	12/2015	Bahadur et al.	N/A	N/A
9843624	12/2016	Taaghoul et al.	N/A	N/A
9866433	12/2017	Fakhouri et al.	N/A	N/A
10523531	12/2018	Zuerner	N/A	N/A
10887276	12/2020	Parulkar et al.	N/A	N/A
11153406	12/2020	Sawant	N/A	H04L 67/561
11188376	12/2020	Alexander et al.	N/A	N/A
11375005	12/2021	Rolando et al.	N/A	N/A
11599376	12/2022	Viswanathan et al.	N/A	N/A
2017/0005983	12/2016	Doukhvalov et al.	N/A	N/A
2018/0367578	12/2017	Verma et al.	N/A	N/A
2019/0109821	12/2018	Clark et al.	N/A	N/A
2020/0052969	12/2019	Xu et al.	N/A	N/A

2020/0067831	12/2019	Spraggins et al.	N/A	N/A
2021/0105225	12/2020	Stammers et al.	N/A	N/A
2021/0152513	12/2020	Grayson et al.	N/A	N/A
2021/0297891	12/2020	Berzin et al.	N/A	N/A
2022/0060408	12/2021	Manickam et al.	N/A	N/A
2022/0103597	12/2021	Gobena et al.	N/A	N/A
2022/0166755	12/2021	Moore et al.	N/A	N/A
2022/0182332	12/2021	Christofi et al.	N/A	N/A
2022/0400053	12/2021	Liu et al.	N/A	N/A
2023/0123775	12/2022	Sivakumar et al.	N/A	N/A

OTHER PUBLICATIONS

Non-Published Commonly Owned Related International Patent Application PCT/US2023/031449 with similar specification, filed Aug. 29, 2023, 55 pages, VMware, Inc. cited by applicant
PCT International Search Report and Written Opinion, PCT/US2023/031449, Nov. 20, 2023, 13 pages. cited by applicant

Non-Published Commonly Owned Related U.S. Appl. No. 18/071,536 with similar specification, filed Nov. 29, 2022, 55 pages, VMware, Inc. cited by applicant

Non-Published Commonly Owned Related U.S. Appl. No. 18/071,537 with similar specification, filed Nov. 29, 2022, 56 pages, VMware, Inc. cited by applicant

Non-Published Commonly Owned Related U.S. Appl. No. 18/071,540 with similar specification, filed Nov. 29, 2022, 56 pages, VMware, Inc. cited by applicant

Non-Published Commonly Owned Related U.S. Appl. No. 18/071,542 with similar specification, filed Nov. 29, 2022, 55 pages, VMware, Inc. cited by applicant

Non-Published Commonly Owned U.S. Appl. No. 18/071,544, filed Nov. 29, 2022, 54 pages, VMware, Inc. cited by applicant

Non-Published Commonly Owned U.S. Appl. No. 18/071,545, filed Nov. 29, 2022, 54 pages, VMware, Inc. cited by applicant

Non-Published Commonly Owned U.S. Appl. No. 18/071,547, filed Nov. 29, 2022, 55 pages, VMware, Inc. cited by applicant

Non-Published Commonly Owned U.S. Appl. No. 18/071,549, filed Nov. 29, 2022, 54 pages, VMware, Inc. cited by applicant

Non-Published Commonly Owned U.S. Appl. No. 18/071,552, filed Nov. 29, 2022, 54 pages, VMware, Inc. cited by applicant

Non-Published Commonly Owned U.S. Appl. No. 18/071,553, filed Nov. 29, 2022, 55 pages, VMware, Inc. cited by applicant

Non-Published Commonly Owned U.S. Appl. No. 18/071,554, filed Nov. 29, 2022, 55 pages, VMware, Inc. cited by applicant

Primary Examiner: Mizrahi; Diane D

Attorney, Agent or Firm: Quarles & Brady LLP

Background/Summary

BACKGROUND

(1) Today, Wi-Fi and other access technologies are utilized for providing private mobile networks deployed as a service for enterprise customers. However, for use cases that have more particular

requirements around factors, such as coverage, reliability, and latency, these access technologies do not suffice. For these scenarios, a private mobile network based on cellular technology (e.g., 4G, 5G, etc.) is a more appropriate solution. Current approaches to such private mobile networks however have not been seamlessly integrated into proven enterprise cloud-native technologies such as Software Defined WAN (SD-WAN), Edge compute and Secure Services Edge (SSE). With an integrated architecture, other challenges emerge such as global management of dispersed network elements, methods to provide system resiliency and end-to-end quality assurance. Finally, novel modifications of the basic architectural approach can be used to cover use cases such as Fixed Wireless Access for rural areas where macro service providers do not operate.

BRIEF SUMMARY

(2) Some embodiments of the invention provide a method implementing a software-defined private mobile network (SD-PMN) for an entity (e.g., a corporation, an educational institution, etc.). To implement the SD-PMN, the method first deploys, at a physical location of the entity, a first set of control plane components for the SD-PMN, including a security gateway, a user-plane function (UPF), an AMF (access and mobility management function), and an SMF (session management function). At an SD-WAN (software-defined wide area network) PoP (point of presence) belonging to a provider of the SD-PMN, the method deploys a second set of control plane components for the SD-PMN, the second set of control plane components comprising a subscriber database that stores data associated with users of the SD-PMN. The method uses an SD-WAN edge router located at the physical location of the entity and a SD-WAN gateway located at the SD-WAN PoP to establish a connection (e.g., via a set of physical network links) from the physical location of the entity to the SD-WAN PoP.

(3) In some embodiments, the security gateway is configured to establish an IPsec (Internet protocol security) tunnel with each access point in a set of access points deployed in the physical location. The access points, in some embodiment, provide a connection between user devices operating in the physical location and the SD-PMN. In some embodiments, the access points receive data message traffic from the user devices as radio waves and convert the radio waves to bits and bytes which are encapsulated and sent to the security gateway via the IPsec tunnels. The security gateway forwards the encapsulated traffic to the UPF, in some embodiments, which decapsulates the traffic and forwards the decapsulated traffic as IP (Internet protocol) traffic to the SD-WAN edge routers.

(4) The connection established by the SD-WAN edge router with the SD-WAN gateway at the SD-WAN PoP is a DMPO (dynamic multipath optimization) tunnel, according to some embodiments, that is established over a set of physical network links (e.g., MPLS, cable modem, 5G, etc.). In some embodiments, the SD-WAN edge router is one of multiple SD-WAN edge routers at multiple physical locations (e.g., branch sites) of the entity, and the SD-WAN gateway is one of multiple SD-WAN gateways at multiple SD-WAN PoPs that each include an instance of the second set of control plane components. Each SD-WAN edge router at each physical location, in some embodiments, is assigned at least a primary SD-WAN gateway associated with a first SD-WAN PoP and a secondary SD-WAN gateway associated with a second SD-WAN PoP for accessing both the internet and components of the SD-WAN PoPs.

(5) In some embodiments, each instance of the control plane shares a same common IP address such that the common IP address can be reached at any of the multiple SD-WAN PoPs. Also, in some embodiments, the SD-PMN is a multi-tenant SD-PMN, and the common IP address associated with the control plane components at the SD-WAN PoPs is the same for all tenants receiving the SD-PMN service. By having a common IP address for the control plane components at each SD-WAN PoP, some embodiments are able to provide SD-WAN resiliency for the SD-PMN, as well as seamless failover between PoPs.

(6) The SD-PMN is centrally managed, in some embodiments, by a private mobile network orchestrator (PMNO). In some such embodiments, the PMNO receives (e.g., from a network

administrator through a user interface (UI) provided by the PMNO) for each physical location in a set of physical locations spanned by the SD-PMN, a tracking area code (TAC) defined for the physical location, as well as a data network name (DNN) defined for each data network in a set of data networks within the SD-PMN. Each data network in the set of data networks span the set of physical locations, according to some embodiments. After receiving the TACs and DNNs, the PMNO of some embodiments stores the TACs and the DNNs in a core of the SD-PMN for use in managing data message traffic and user devices in the SD-PMN.

(7) Some embodiments of the invention also provide a method for deploying a private mobile network for an entity in particular geographic area. For each physical location in a set of physical locations within the particular geographic area (e.g., homes in a municipality), the method deploys an SD-WAN (software-defined wide area network) customer premise appliance, such as an SD-WAN edge router enabled with LTE, 4G, or 5G. The method establishes, via a RAN (radio access network), a set of connections between each SD-WAN customer premise appliance and at least one physical access point in a set of physical access points deployed in the particular geographic area. Each physical access point has a connection to a central aggregation point that includes a UPF (user plane function) and an SD-WAN edge router that connect the physical access nodes to an SD-WAN PoP (point of presence). The method then uses the established set of connections to provide a private mobile networking service for the particular geographic area.

(8) The preceding Summary is intended to serve as a brief introduction to some embodiments of the invention. It is not meant to be an introduction or overview of all inventive subject matter disclosed in this document. The Detailed Description that follows and the Drawings that are referred to in the Detailed Description will further describe the embodiments described in the Summary as well as other embodiments. Accordingly, to understand all the embodiments described by this document, a full review of the Summary, the Detailed Description, the Drawings, and the Claims is needed. Moreover, the claimed subject matters are not to be limited by the illustrative details in the Summary, the Detailed Description, and the Drawings.

Description

BRIEF DESCRIPTION OF FIGURES

- (1) The novel features of the invention are set forth in the appended claims. However, for purposes of explanation, several embodiments of the invention are set forth in the following figures.
- (2) FIG. 1 conceptually illustrates a first example of an architecture diagram of a multi-tenant PMN of some embodiments.
- (3) FIG. 2 conceptually illustrates a process of some embodiments performed to implement an SD-PMN such as the SD-PMN described above for FIG. 1.
- (4) FIG. 3 conceptually illustrates a SD-PMN of some embodiments that includes multiple branch sites and multiple SD-WAN PoPs.
- (5) FIG. 4 conceptually illustrates an architecture diagram showing a data plane flow through an SD-PMN of some embodiments.
- (6) FIG. 5 conceptually illustrates an architecture diagram showing a control plane flow through an SD-PMN of some embodiments.
- (7) FIG. 6 conceptually illustrates a second example of an architecture diagram of a multi-tenant PMN of some embodiments.
- (8) FIG. 7 conceptually illustrates a process of some embodiments for establishing an SD-PMN such as the SD-PMN illustrated in the FIG. 6.
- (9) FIG. 8 conceptually illustrates an architecture diagram in which a multi-tenant SD-PMN is implemented across customer premises of some embodiments that have a disaggregated RAN.
- (10) FIG. 9 conceptually illustrates a failover process performed by an SD-WAN edge router of

some embodiments when connecting to an SD-WAN PoP.

(11) FIG. 10, which conceptually illustrates a set of diagrams showing failover from a first PoP to a second PoP.

(12) FIG. 11 conceptually illustrates a diagram of a SD-PMN having a centralized management plane, in some embodiments.

(13) FIG. 12 conceptually illustrates a diagram of the OAM architecture of some embodiments for a multi-tenant SD-PMN.

(14) FIG. 13 conceptually illustrates an architecture diagram of a centrally managed SD-PMN of some embodiments.

(15) FIG. 14 conceptually illustrates a process of some embodiments for centrally managing a SD-PMN.

(16) FIG. 15 conceptually illustrates a process of some embodiments for implementing an SD-PMN as a fixed wireless network for a particular geographic location.

(17) FIG. 16 conceptually illustrates the architecture of a fixed wireless network of some embodiments.

(18) FIG. 17 conceptually illustrates a diagram of a joint orchestration platform of some embodiments that orchestrates applications end-to-end across an SD-PMN, SD-WAN, and edge compute stack.

(19) FIG. 18 illustrates a joint orchestration platform diagram that includes a set of example intent-based APIs for orchestrating a video analytics application, in some embodiments, across an SD-PMN, SD-WAN, and edge compute stack.

(20) FIG. 19 conceptually illustrates a computer system with which some embodiments of the invention are implemented.

DETAILED DESCRIPTION

(21) In the following detailed description of the invention, numerous details, examples, and embodiments of the invention are set forth and described. However, it will be clear and apparent to one skilled in the art that the invention is not limited to the embodiments set forth and that the invention may be practiced without some of the specific details and examples discussed.

(22) Some embodiments of the invention provide a method implementing a software-defined private mobile network (SD-PMN) for an entity (e.g., a corporation, an educational institution, etc.). To implement the SD-PMN, the method first deploys, at a physical location of the entity, a first set of control plane components for the SD-PMN, including a security gateway, a user-plane function (UPF), an AMF (access and mobility management function), and an SMF (session management function). At an SD-WAN (software-defined wide area network) PoP (point of presence) belonging to a provider of the SD-PMN, the method deploys a second set of control plane components for the SD-PMN, the second set of control plane components comprising a subscriber database that stores data associated with users of the SD-PMN. The method uses an SD-WAN edge router located at the physical location of the entity and a SD-WAN gateway located at the SD-WAN PoP to establish a connection from the physical location of the entity to the SD-WAN PoP.

(23) In some embodiments, the security gateway is configured to establish an IPsec (Internet protocol security) tunnel with each access point in a set of access points deployed in the physical location. The access points, in some embodiment, provide a connection between user devices operating in the physical location and the SD-PMN. In some embodiments, the access points receive data message traffic from the user devices as radio waves and convert the radio waves to bits and bytes which are encapsulated and sent to the security gateway via the IPsec tunnels. The security gateway forwards the encapsulated traffic to the UPF, in some embodiments, which decapsulates the traffic and forwards the decapsulated traffic as IP (Internet protocol) traffic to the SD-WAN edge routers.

(24) The connection established by the SD-WAN edge router with the SD-WAN gateway at the

SD-WAN PoP is a DMPO (dynamic multipath optimization) tunnel, according to some embodiments. In some embodiments, the SD-WAN edge router is one of multiple SD-WAN edge routers at multiple physical locations (e.g., branch sites) of the entity, and the SD-WAN gateway is one of multiple SD-WAN gateways at multiple SD-WAN PoPs that each include an instance of the second set of control plane components. Each SD-WAN edge router at each physical location, in some embodiments, is assigned at least a primary SD-WAN gateway associated with a first SD-WAN PoP and a secondary SD-WAN gateway associated with a second SD-WAN PoP for accessing both the internet and components of the SD-WAN PoPs.

(25) In some embodiments, each instance of the control plane shares a same common IP address such that the common IP address can be reached at any of the multiple SD-WAN PoPs. Also, in some embodiments, the SD-PMN is a multi-tenant SD-PMN, and the common IP address associated with the control plane components at the SD-WAN PoPs is the same for all tenants receiving the SD-PMN service. By having a common IP address for the control plane components at each SD-WAN PoP, some embodiments are able to provide SD-WAN resiliency for the SD-PMN, as well as seamless failover between PoPs.

(26) The SD-PMN is centrally managed, in some embodiments, by a private mobile network orchestrator (PMNO). In some such embodiments, the PMNO receives (e.g., from a network administrator through a user interface (UI) provided by the PMNO) for each physical location in a set of physical locations spanned by the SD-PMN, a tracking area code (TAC) defined for the physical location, as well as a data network name (DNN) defined for each data network in a set of data networks within the SD-PMN. Each data network in the set of data networks span the set of physical locations, according to some embodiments. After receiving the TACs and DNNs, the PMNO of some embodiments stores the TACs and the DNNs in a core of the SD-PMN for use in managing data message traffic and user devices in the SD-PMN.

(27) Some embodiments of the invention also provide a method for deploying a private mobile network for an entity in particular geographic area. For each physical location in a set of physical locations within the particular geographic area (e.g., homes in a municipality), the method deploys an SD-WAN (software-defined wide area network) customer premise appliance, such as an SD-WAN edge router enabled with LTE, 4G, or 5G. The method establishes, via a RAN (radio access network), a set of connections between each SD-WAN customer premise appliance and at least one physical access point in a set of physical access points deployed in the particular geographic area. Each physical access point has a connection to a central aggregation point that includes a UPF (user plane function) and an SD-WAN edge router that connect the physical access nodes to an SD-WAN PoP (point of presence). The method then uses the established set of connections to provide a private mobile networking service for the particular geographic area.

(28) FIG. 1 conceptually illustrates a first example of an architecture diagram of a multi-tenant PMN of some embodiments. In this architecture diagram **100**, components of the control plane are split between locations on customer premises **101** and **102** and the SD-WAN PoP **105**. For instance, in this particular architecture **100**, each customer premises **101** and **102** includes a security gateway **124**, a UPF **126**, and an AMF/SMF **128**, while the SD-WAN PoP **105** includes the multi-tenant 5G unified data management (UDM) for storing user data. It should be noted that the components illustrated in the architecture in FIG. 1 are components associated with a 5G network, and other embodiments of the invention can include components associated with other network types, such as MMEs (mobility management entities) for 4G and LTE solutions.

(29) The user devices **130** and **135**, in some embodiments, can include mobile telephones, tablets, computers (e.g., desktop computers, laptop computers, etc.), and any other devices belonging to users of the PMN that have the ability to connect to the internet. The access nodes **122** and **152** in the local RAN's **120** and **150** receive control plane and data plane data message traffic from the user device **130-135**. The access nodes **122** and **152** are physical access points (i.e., base stations, radio units, etc.) that are part of the RAN, in some embodiments, and provide user devices with

access to the PMN. For instance, for LTE (long term evolution) networks, the access nodes **122** and **152** are eNodeBs (eNBs), while for 5G networks, the access nodes **122** and **152** are gNBs (next generation NodeBs).

(30) In some embodiments, the data message traffic is transmitted from the user devices **130** and **135** to the access nodes **122** and **152** via radio signals. Upon receiving the data message traffic from the user devices **130-135**, the access nodes **122** and **152** process the data message traffic, convert the radio signals into collections of bits and bytes, encapsulate the data message traffic (i.e., the bits and bytes) and forward the encapsulated traffic to the security gateways **124** and **154** via tunnels between the access nodes **122** and **152** and the security gateways **124** and **154**.

(31) The tunnels through which the access nodes **122** and **152** forward data message traffic to the security gateways **124** and **154** are IPsec tunnels established by the security gateways **124** and **154**, according to some embodiments. In some embodiments, the data message traffic is encapsulated as GPRS (general packet radio service) tunneling protocol (GTP) before being forwarded through the IPsec tunnels between the access nodes and security gateways. The security gateways **124** and **154** decapsulate the received encapsulated data message traffic, and depending on the traffic type (i.e., control traffic or dataplane traffic), forward the data message traffic as IP traffic to either their respective UPFs **126** and **156**, or their respective AMFs/SMFs **128** and **158**. While illustrated as a single element, it should be noted that the AMF and SMF, as will be described below, are separate functions that perform different functionalities of the control plane and in other embodiments of the invention may be illustrated separately. The UPFs **126** and **156** handle data plane traffic, according to some embodiments, while the AMFs/SMFs **128** and **158** are responsible for control plane traffic. Additional details regarding the UPFs, AMFs, and SMFs will be provided below by reference to FIGS. 4-5.

(32) In order to send traffic to destinations external to the customer premises **101** and **102**, and, in some embodiments, to certain destinations within the customer premises **101** and **102**, the UPFs **126** and **156** and the AMFs/SMFs **128** and **158** forward the data message traffic to the SD-WAN edge routers **140** and **145**. The SD-WAN edge routers **140** and **145** each connect to the SD-WAN gateway **110** to provide connections between the customer premises **101** and **102** and the SD-WAN PoP **105**, as shown. The SD-WAN edge routers **140** and **145** of some embodiments establish DMPO tunnels to the SD-WAN gateway **110**, as well as to other SD-WAN gateways (not shown), and, in some embodiments, with each other and other SD-WAN components (e.g., other SD-WAN edge routers at other physical locations of the entity, and SD-WAN hubs).

(33) In some embodiments, once an SD-WAN edge router **140-145** detects a WAN (wide area network) link (e.g., a WAN link between the SD-WAN edge router and an SD-WAN gateway), the SD-WAN edge router **140-145** establishes a DMPO tunnel with the SD-WAN gateway **110** and run bandwidth tests using short bursts of bi-directional traffic with the SD-WAN gateway **110** (or a different SD-WAN gateway (not shown) that is closer in proximity to the SD-WAN edge router than the SD-WAN gateway **110**). Because the SD-WAN gateway **110** is deployed at the SD-WAN PoP **105**, the SD-WAN gateway **110** can identify real public IP addresses of WAN links in cases where a NAT (network address translation) or PAT (port address translation) device sits in front of the SD-WAN edge router, or where the WAN link is a private link, according to some embodiments.

(34) Once the DMPO tunnels are established, in some embodiments, DMPO performs uni-directional performance metric measurements for metrics such as loss, latency, and jitter, for every data message on every DMPO tunnel between two DMPO endpoints (i.e., the SD-WAN edge routers **140-145** and the SD-WAN gateway **110**). In some embodiments, the DMPO tunnel header used to encapsulate each packet sent through the DMPO tunnel includes performance metrics such as sequence number and timestamp to enable the DMPO endpoints to identify lost packets and out-of-order packets, as well as to compute jitter and latency bi-directionally. These performance metrics are communicated between the DMPO endpoints at an order of every 100 ms, according to

some embodiments, and when there is no active data message traffic being sent through the DMPO tunnels, the DMPO endpoints instead send active probes every 100 ms, or every 500 ms after a certain period (e.g., 5 minutes) of no high-priority data message traffic. Additionally, for any private WAN links for which a class of service (CoS) agreement is defined, DPMO is configured in some embodiments to take the CoS agreement into account for all traffic steering decisions (e.g., monitoring traffic, data plane application traffic, etc.). DMPO can also add Forward Error Correction (FEC) for certain classes of traffic, according to some embodiments.

(35) Once the SD-WAN gateway **110** receives data message traffic from the SD-WAN edge routers **140** and **145**, the SD-WAN gateway **110** determines whether the data message traffic is application traffic or control plane traffic, and either forwards the data message traffic to the internet or to the UDM **112**. In some embodiments, before forwarding internet traffic to the internet, the SD-WAN gateway **110** forwards the traffic to the SASE services **114** for processing. The SASE services **114**, in some embodiments, are provided by a third-party vendor and can include services such as firewall as a service, secure web gateway, zero-trust network access, and other threat detection services.

(36) FIG. 2 conceptually illustrates a process of some embodiments performed to implement an SD-PMN such as the SD-PMN described above for FIG. 1. The process **200** starts when the process deploys (at **210**) a security gateway, UPF, AMF, and SMF to each physical location in a set of physical locations across which the SD-PMN is being deployed. Each physical location of the entity also includes a SD-WAN edge router deployed to the physical location in order to provide a connection between the physical location and SD-WAN PoPs of the provider of the SD-PMN, as well as between the physical location and other physical locations and hub datacenters of the entity. As mentioned above, the SD-WAN edge routers and SD-WAN gateways utilize SD-WAN services such as DMPO, in some embodiments, to optimize the data message traffic sent between SD-WAN devices implemented in the SD-PMN.

(37) The process **200** deploys (at **220**) a subscriber database (i.e., UDM) that stores data associated with users of the PMN that belong to the particular entity for which the PMN is being deployed to each of a set of SD-WAN PoPs belonging to a provider of the PMN. For instance, in the architecture diagram **100** described above, the UDM **112** is located in the PoP **105** while the security gateways **124** and **154**, UPFs **126** and **156**, and AMFs/SMFs **128** and **158** are located on the customer premises **101** and **102**. The user data stored by the UDM, in some embodiments, includes customer profile information, customer authentication information, and, in some embodiments, a set of encryption keys for the information. As the UDM (i.e., instances of the UDM) is deployed to each PoP, the SD-WAN edge routers can connect to any SD-WAN gateway for any of the SD-WAN PoPs.

(38) For each SD-WAN edge router deployed at each physical location in the set of physical locations across which the SD-PMN is being deployed, the process **200** provides a list of SD-WAN gateways to which the SD-WAN edge router can connect to access elements of the SD-WAN PoP. For example, FIG. 3 conceptually illustrates a SD-PMN **300** that includes multiple branch sites **310**, **312**, and **314** and multiple SD-WAN PoPs **320**, **322**, and **324**. Each branch site **310-314** includes a respective SD-WAN edge router **330**, **332**, and **324** and set of machines **350**, **352**, and **354** (e.g., user devices and other network devices deployed to the branch sites), as shown. Each of the SD-WAN PoPs **320-324** includes a respective SD-WAN gateway **340**, **342**, and **344**, and a set of resources **360**, which are the same for each SD-WAN PoP.

(39) Each of the branch sites **310-314** includes a respective gateway list **370**, **372**, and **374** provided to the SD-WAN edge routers **330-334** for use in connecting their respective branch sites **310-314** to resources **360** located in the PoPs **320-324**, with each gateway list **370-374** including a primary gateway and a secondary gateway. Accordingly, each of the SD-WAN edge routers **330-334** has established a connection (e.g., DMPO tunnel) **380** to their primary SD-WAN gateways and a connection **385** to their secondary SD-WAN gateways. As will be described in the embodiments

further below, the secondary gateways are for use by the SD-WAN edge routers 330-334 for failover. Returning to the process 200, after the gateway lists have been provided to the SD-WAN edge routers, the process 200 ends.

(40) As mentioned above regarding the architecture diagram 100 of FIG. 1, the UPFs 126 and 156 handle data plane traffic, while the AMFs/SMFs 128 and 158 are responsible for control plane traffic. FIG. 4 conceptually illustrates an architecture diagram showing a data plane flow through an SD-PMN of some embodiments. The diagram 400 includes a customer premises 401 and an SD-WAN PoP 405. The customer premises 401 includes user devices 430, a local RAN 420 that includes at least one access node 422, a security gateway 424, a UPF 426, an AMF/SMF 428, an SD-WAN edge 440, and on-premise destinations 450. The SD-WAN PoP 405 includes an SD-WAN gateway 410, a multi-tenant 5G UDM 412, and SASE services 414.

(41) For user devices 430 that are already authenticated with the SD-PMN, the data message traffic sent from these devices is transmitted via radio waves from the user devices 430 to the access nodes 422 that are part of the local RAN 420. The access node 422 processes the received data message traffic radio waves and converts the radio waves into bits and bytes, and encapsulates the data to be transmitted to the security gateway 424 via an IPsec tunnel established by the security gateway 424 between the access node 422 and security gateway 424. When the data message traffic (i.e., the encapsulated bits and bytes) arrives at the security gateway 424, the security gateway 424 decapsulates the traffic and forwards the decapsulated traffic to the UPF 426. In some embodiments, the decapsulated traffic is GTP traffic.

(42) When the UPF 426 receives the GTP traffic from the security gateway 424, the UPF 426 then removes the GTP header from the data message traffic and sends the data message traffic out as IP traffic. In some embodiments, data messages having destinations within the customer premises 401 are sent directly from the UPF 426 to their intended destinations. In other embodiments, the UPF 426 sends the data message traffic to the SD-WAN edge router 440, which then forwards the data message traffic to, e.g., an on-premise destination 450. Similarly, for data message traffic associated with destinations external to the customer premise 401, such as the SD-WAN PoP 405, the UPF 426 forwards the data message traffic to the SD-WAN edge router 440, which forwards the traffic to a next-hop forwarding elements, such as the SD-WAN gateway 110 through a DMPO tunnel established between these endpoints, which provides access to the SD-WAN PoP 405, and also forwards traffic having destinations external to the SD-PMN, such as the internet 460.

(43) The SD-WAN gateway 410, in some embodiments, determines whether the received data message traffic requires processing by, e.g., the SASE services 414, or any other packet processing pipelines deployed to the SD-WAN PoP 405 (e.g., middlebox service engines). For data message traffic that does require processing by the SASE services 414, or other processing performed within the SD-WAN PoP 405, the SD-WAN gateway 410 forwards the data message traffic for processing, and then subsequently forwards the processed data message traffic to its destination, such as via the internet 460.

(44) FIG. 5 conceptually illustrates an architecture diagram showing a control plane flow through an SD-PMN of some embodiments. Like the diagram 400, the diagram 500 includes a customer premises 501 and an SD-WAN PoP 505. The customer premises 501 includes user devices 530, a local RAN 520 that includes at least one access node 522, a security gateway 524, a UPF 526, an AMF 528, an SMF 528b, and an SD-WAN edge 540. The SD-WAN PoP 505 includes an SD-WAN gateway 510, a multi-tenant 5G UDM 512, and SASE services 514.

(45) When a user device 530 attempts to connect to the SD-PMN, the user device 530 tries to associate with an access node 522 that is part of the local RAN 520 to request to join the network, according to some embodiments. The access node 522 then communicates with the security gateway 524 with the intention of the communications reaching the AMF 528a. The AMF 528a is responsible for authenticating the user device 530. In some embodiments, the AMF 528a performs the user authentication by performing a look-up in a database that stores user information to

determine who is trying to join the network, as well as what type of service should be provided to the user device **530** upon authentication.

(46) Accordingly, the AMF **528a** sends a control message (e.g., an authentication request) destined to an IP address associated with the UDM **512** to the SD-WAN edge router **540**, which sends the control messages via a DMPO tunnel to the SD-WAN gateway **510** at the SD-WAN PoP **505** that hosts the UDM **512**. The IP address associated with the UDM **512** is a common IP address at every single SD-WAN PoP connected by the SD-PMN, according to some embodiments. Additional details regarding the use of a common IP address for each SD-WAN PoP will be described in further detail by reference to embodiments below.

(47) Upon receiving the authentication request, the UDM **512** performs user authentication by identifying in its database which user is trying to join based on an identifier associated with the user, whether the user is authorized to join the network, as well as the kind of service should be provided to the user if the user is authenticated. For user devices that are not authorized, the UDM **512** indicates to the AMF **528a** that the device is not authorized, and the AMF **528a** subsequently provides that response back to the unauthorized user device. Otherwise, when the UDM **512** determines that the user is allowed to join the network, the UDM **512** sends a reply to the AMF **528a** that indicates that the user device **530** is authorized to use the network and provides any necessary information about the user and/or user device.

(48) In response to a user device **530** being authorized, the SMF **528b** then communicates with the UPF **526** to direct the UPF **526** to set up a bearer for the newly authorized user device **530**. The UPF **526** then provides the user device with an IP address, enables a particular QoS (quality of service) for the user device's communications, and indicates which subnets (e.g., VLANs) to put the user device's traffic on. In some embodiments, when applicable, the SMF **528b** directs the UPF **526** to also create certain network slices for the newly authorized user device. Once the UPF **526** has completed these steps, the user device **530** can communicate using the SD-PMN.

(49) In some embodiments, a different control plane architecture is utilized for implementing an SD-PMN. FIG. 6 conceptually illustrates a second example of an architecture diagram of a multi-tenant PMN of some embodiments. The architecture diagram **600** includes customer premises **601** and **602** for first and second customers respectively, as well as an SD-WAN PoP **605**. Each customer premises **601** and **602** includes respective user devices **630** and **635**, a local RAN **620** and **650** that includes at least one access node **622** and **652**, a UPF **624** and **654**, and an SD-WAN edge router **640** and **645**. The SD-WAN PoP **605** includes an SD-WAN gateway **610**, SASE services **618**, a security gateway **612**, and a multi-tenant control plane **616** that includes an AMF **660**, SMF **665**, and UDM **614**. That is, unlike the embodiments described above, the security gateway, AMF, and SMF are deployed to the PoPs rather than to each customer premise. The architecture diagram **600** will be described in more detail below by reference to FIG. 7.

(50) FIG. 7 conceptually illustrates a process of some embodiments for establishing an SD-PMN such as the SD-PMN illustrated in the diagram **600**. The process **700** starts when the process deploys (at **710**) a UPF to each physical location in a set of physical locations belonging to a particular entity for which the SD-PMN is being established. The UPF **624** and **654** operates as an interconnecting point between the local RAN **620** and **650** and components of the data network (e.g., the SD-WAN edge router **640** and **645**), performs routing and forwarding for data messages sent on the control plane and data plane, and performs various other functions such as setting up bearers for newly authorized user devices. For instance, because the AMF **660** and SMF **665** are deployed to the SD-WAN PoP **605** instead of to each customer premise **601** and **602**, the UDF **624** and **654** receives and forwards control data messages, such as requests from new user devices looking to access the SD-PMN.

(51) When a new user device attempts to access the SD-PMN, the request is received by the access point **622** and **652** that is part of the local RAN **620** and **650**, which converts the request from radio waves to bits and bytes, encapsulates the converted request, and forwards the encapsulated request

via a GTP tunnel to the UPF **624** and **654**. The UPF **624** and **654** sends the request as IP traffic to the SD-WAN edge router **640** and **645**, which forwards the request via a DMPO tunnel between the SD-WAN edge router **640** and **645** and the SD-WAN gateway **610**. The SD-WAN gateway **610** forwards the request to the security gateway **612**, which provides the request to the multi-tenant 5G control plane **616**.

(52) Returning to the process **700**, the process deploys (at **720**) a security gateway, AMF, SMF, and subscriber database that stores data associated with users of the SD-PMN that belong to the particular entity (an any other entity that uses the multi-tenant SD-PMN) to each of a set of SD-WAN PoPs belonging to a provider of the SD-PMN. That is, rather than just deploying the UDM **614** to the SD-WAN PoPs **605**, the AMF **660** and SMF **665** are also deployed to the SD-WAN PoPs **605**. As such, to continue to process an authentication request for a new user, the security gateway **612** provides the request to the AMF **660** that is deployed to the multi-tenant 5G control plane **616**.

(53) Upon received an authentication request, the AMF **660** performs a lookup in the UDM **614** to determine whether the requesting device is allowed to access the SD-PMN. For requesting devices that are not authorized, the AMF **660** sends a response back to the UPF **624** and **654** via the security gateway **612**, SD-WAN gateway **610**, and SD-WAN edge **640** and **645**, which then notifies (i.e., via the local RAN **620** and **650**) that the device is not authorized to access the SD-PMN. For devices that are authorized, the SMF **665** then communicates with the UPF **624** and **654** to direct the UPF to set up a bearer for the newly authorized and authenticated user device by providing the user device with an IP address, enabling a particular QoS for the user device's communications, and indicating which subnet(s) to put the user device's traffic on.

(54) For each SD-WAN edge device deployed at each physical location in the set of physical locations, the process **700** provides (at **730**) a list of SD-WAN gateways that provide access to the SD-WAN PoPs to which the SD-WAN edge router can connect to access the SD-WAN PoPs. While only one SD-WAN gateway **610** at one SD-WAN PoP **605** is illustrated in the diagram **600**, other embodiments of the invention include multiple SD-WAN PoPs each having a respective SD-WAN gateway. Each SD-WAN edge router at each branch site (i.e., customer premise) then receives a list indicating at least a primary and secondary SD-WAN gateway for the SD-WAN edge router to connect to in order to access services and components of the SD-WAN PoPs, such as the control plane deployed to the PoPs. Following **730**, the process **700** ends.

(55) The distributed control plane architecture described above for FIG. **6** is applicable for both small cell and disaggregated RANs, according to some embodiments. FIG. **8** conceptually illustrates an architecture diagram in which a multi-tenant SD-PMN is implemented across customer premises of some embodiments that have a disaggregated RAN. As shown, the diagram **800** includes two customer premises **801** and **802** and an SD-WAN PoP **805**. The SD-WAN PoP **805** includes an SD-WAN gateway **810**, security gateway **812**, multi-tenant 5G control plane **816**, and SASE services **818**. Similar to the diagram **600**, the multi-tenant 5G control plane **816** includes a UDM **814**, an AMF **860**, and an SMF **865**. Each customer premises **801** and **802** includes a respective local RAN **820** having at least one access point, a distributed unit (DU) **822** and **852**, a central unit (CU) **824** and **854**, a configuration server **826** and **856**, X86 microprocessors **828** and **858**, a UPF **860** and **865**, and an SD-WAN edge router **840** and **845**.

(56) On each customer premise **801** and **802**, the local RANs **820** and **850** act as the radio units (RUs) for the disaggregated RAN and enable geographical coverage using radio functions. The DUs **822** and **852** realize baseband processing functions across the different physical locations spanned by the SD-PMN as virtualized network functions that run on hardware, and allow for possible hardware acceleration. The CUs **824** and **854** centralize data message processing functions, and, like the DUs, realize these functions as virtualized network functions that run on hardware. The configuration servers **826** and **856** receive configuration data for the disaggregated RAN from a manager (e.g., management server) for the SD-PMN.

(57) When a new user device attempts to access the SD-PMN, the user device communicates via

radio waves with the local RAN **820** and **850**, which converts the radio waves to bits and bytes, encapsulates the bits and bytes, and transmits the encapsulated bits and bytes through a tunnel to the UPF **860** and **865**. The UPF **860** and **865** then transmit the authentication request to the SD-WAN edge router **840** and **845** as IP traffic, and the SD-WAN edge router **840** and **845** uses the DMPO tunnel to the SD-WAN gateway **810** to forward the request to the SD-WAN PoP **805**.

(58) The AMF **870** performs a look-up with the UDM **814** to determine whether the requesting user device is authorized to use the SD-PMN. For devices that are not authorized, the AMF **870** sends a response back toward the UPF **860** to indicate the requesting device is not authorized, and the UPF **860** notifies the requesting device that the request is denied. For devices that are authorized, the SMF **875** then communicates with the UPF **860** to have the UPF **860** set up a bearer for the new user device to enable the new user device to use the SD-PMN. The UPF **860** and **865** then provides the user device with an IP address, enables a particular QoS (quality of service) for the user device's communications, and indicates which subnets (e.g., VLANs) to put the user device's traffic on.

(59) It should be noted that in the embodiments described above and below that for the core network, the data plane components are located on customer premises, the control plane components are split between customer premises and the SD-WAN PoPs, and, as will be described below, the management plane components are centrally located in the cloud. Additionally, for the RAN, the data plane components and control plane components are located on customer premises, and the management plane components are centrally located in the cloud (or multiple clouds such as in the case of multiple RAN vendors as will be described further below). As will be described in some embodiments below, the core network management plane components and the RAN management plane components of some embodiments are located in separate clouds that connect to a unified and centralized management system for the SD-PMN.

(60) In some embodiments, the IP address assigned to the control plane components (i.e., the control plane components for the core network) deployed to the SD-WAN PoPs is a common IP address that is used for those components across all of the PoPs, regardless of location or customer. In other words, all of the control plane components for the multi-tenant SD-PMN are reachable at the same IP address in each SD-WAN PoP. As a result, the SD-PMN is able to provide resiliency and seamless failover between SD-WAN PoPs. FIG. **9** conceptually illustrates a failover process performed by an SD-WAN edge router of some embodiments when connecting to an SD-WAN PoP. The process **900** will be described in detail below with references to FIG. **10**, which conceptually illustrates a set of diagrams showing failover from a first PoP to a second PoP.

(61) The process **900** starts when the SD-WAN edge router receives (at **910**) a list of SD-WAN gateways associated with SD-WAN PoPs of the SD-PMN to which the SD-WAN edge router can connect to access an SD-WAN PoP. In some embodiments, when the SD-PMN is established, each SD-WAN edge router at each branch location of an entity for which the SD-PMN is implemented receives a gateway list that specifies at least a primary SD-WAN gateway and secondary SD-WAN gateway to which the SD-WAN edge router can connect to access an SD-WAN PoP. In the diagrams **1001** and **1002**, for instance, the SD-WAN edge router **1005** at the branch site **1015** is illustrated as having a gateway list **1010** specifying gateway **1020** as a primary gateway and gateway **1025** as a secondary gateway. In some embodiments, the gateway lists are determined based on proximity of the gateways to the SD-WAN edge router (i.e., the primary gateway being the closest gateway to the edge router).

(62) The process **900** establishes (at **920**) SD-WAN tunnels with primary and secondary SD-WAN gateways identified in the received list. In the diagram **1001**, the SD-WAN edge router **1005** has a first tunnel **1030** to its primary SD-WAN gateway **1020** and a second tunnel **1035** to its secondary SD-WAN gateway **1025** (drawn with a dashed line to indicate it is the tunnel to the secondary gateway).

(63) The process **900** uses (at **930**) the SD-WAN tunnel established with the primary SD-WAN

gateway to forward data messages addressed to a particular IP address associated with the control plane components (i.e., core) located in the SD-WAN PoP. As illustrated in the diagrams **1001** and **1002**, each both the first PoP **1040** and the second PoP **1045** include two active instances of the core **1050**. Each instance of the core **1050** is identical to each other instance of the core **1050**, is stateless, and is reachable at the same common IP address, as described in the embodiments above. As such, in some embodiments, if one active core in a PoP becomes unavailable, the SD-WAN gateway for that PoP can automatically switch to the other active core for forwarding data messages associated with the common IP address assigned to the core, thereby providing resiliency and seamless failover within each PoP.

(64) The process **900** determines (at **940**) that the SD-WAN tunnel between the SD-WAN edge router at the branch location and the primary SD-WAN gateway is down. In the diagram **1002**, for instance, the tunnel **1030** from the SD-WAN edge router **1005** to the SD-WAN gateway is indicated as being down. The SD-WAN tunnel between the SD-WAN edge router and the SD-WAN gateway is a DMPO tunnel, according to some embodiments. Accordingly, in some embodiments, DMPO detects the outage, or brownout (e.g., when a particular link cannot provide SLA for a particular application), while performing its continuous monitoring, and can alert the SD-WAN edge router of the tunnel failure to cause the SD-WAN edge router to use the secondary SD-WAN gateway for reaching a particular IP address in an SD-WAN PoP.

(65) Based on the determination that the SD-WAN tunnel to the primary SD-WAN gateway is down, the process **900** uses (at **950**) the SD-WAN tunnel established with the secondary SD-WAN gateway to forward the data message to the destination IP address located the second SD-WAN PoP. As illustrated in the diagram **1002**, the tunnel **1035** is now shown as a solid line to indicate this tunnel is now the active tunnel being used by the SD-WAN edge router **1005** to connect its branch site **1015** to the core **1050**. In some embodiments, the SD-WAN edge router **1005** continues to use the tunnel **1035** to the secondary SD-WAN gateway **1025** at the second PoP **1045** to reach the core **1050** until the tunnel **1030** is back up and running. In other embodiments, the SD-WAN edge router **1005** uses the tunnel **1035** until the tunnel **1035** experiences an outage or brownout. Following **950**, the process **900** ends.

(66) In some embodiments, the SD-PMN is managed and configured by a centralized management plane. FIG. **11** conceptually illustrates a diagram of a SD-PMN having a centralized management plane, in some embodiments. As shown, the diagram **1100** includes a central orchestrator **1110** located in a cloud **1105**, one or more RAN management OAM (operations, administration, and maintenance) servers **1114** located in a cloud **1112**, a 5G core OAM server **1118** located in a cloud **1116**, an on-premise enterprise location **1120**, an SD-WAN PoP **1122**, and clouds **1124**. The on-premise enterprise location **1120** includes endpoints **1130** with SIM-enabled modems, a small cell RAN **1132**, edge applications **1134**, an SD-WAN edge router **1136**, and a 4G/5G packet core **1138**. The SD-WAN PoP **1122** includes SASE services **1140**, an SD-WAN gateway **1142**, and a 5G core **1144** that includes the UDM for the SD-PMN. The clouds **1124** include a public cloud **1150** that hosts multiple applications **1152**, and a private cloud datacenter **1154** that includes an SD-WAN edge router **1156** for connecting to the SD-WAN PoP **1122** and other elements external to the private datacenter **1154**.

(67) Because the components of the RAN and core are distributed, in some embodiments, with the RAN on each customer's premises and the core split between each customers' premises and the SD-WAN PoPs of the SD-PMN provider, information associated with each portion of the RAN and core needs to be centralized into one management layer in order for, e.g., certain parameters to be assigned. Accordingly, the PMNO **1110**, in some embodiments, centrally stores tenant-specific 5G core and access point (i.e., RAN) information and pushes this information to respective OAM servers **1114** and **1118** as configuration updates are made. In some embodiments, the OAM servers **1114** and **1118** are deployed across multiple administrative domains in one or more public clouds. Also, in some embodiments, the PMNO **1110** is configured to support multiple RAN vendors,

including multiple RAN vendors for a single customer premise. In some such embodiments, the PMNO **1110** invokes vendor-specific APIs against the corresponding RAN OAM server based on both the type of access point associated and the site at which that access point is deployed.

(68) In some embodiments, all sessions between the on-premise components, such as the RAN **1132** and 4G/5G packet core **1138**, and their respective OAM servers in the cloud(s), such as the RAN OAM server(s) **1114** located in the cloud **1112** and 5G core server **1118** located in the cloud **1116**, are initiated by the on-premise components rather than by the cloud components. That is, the cloud components will not establish new connections, according to some embodiments.

Additionally, the connection **1160** between the RAN **1132** and RAN OAM server(s) **1114**, and the connection **1165** between the 4G/5G packet core **1138** and 5G core OAM **1118** are encrypted connections, according to some embodiments.

(69) FIG. **12** conceptually illustrates a diagram of the OAM architecture of some embodiments for a multi-tenant SD-PMN. The diagram **1200** includes a PMNO **1205**, RAN OAM servers **1210a** and **1210b** for a first vendor, RAN OAM servers **1212a** and **1212b** for a second vendor, core OAM servers **1214a** and **1214b**, a load balancer **1220** for connecting to an access point **1230** of the first RAN vendor on a first customer's premise, a load balancer **1222** for connecting to an access point **1234** of the second RAN vendor on a second customer's premise, and a load balancer **1224** for connecting to the distributed 4G/5G packet core components including the 4G/5G packet core components **1232** on the first customer's premise, the 4G/5G packet core components **1236** on the second customer's premise, and the 4G/5G UDM **1238** located in the SASE PoP of the SD-PMN provider (e.g., an SD-WAN PoP). Each OAM server pair has full redundancy as illustrated and are configured as active/active pairs, according to some embodiments.

(70) It should be noted that in some embodiments, rather than terminating at the load balancers **1220** and **1222**, the connections from the access points **1230** and **1234** instead terminate at security gateways associated with the respective RANs. The security gateways in some such embodiments provide unique identifiers for each access point **1230** and **1234** to their respective RAN OAM servers **1210a-1210b** and **1212a-1212b** so that each customer can be uniquely identified through the security gateway.

(71) As mentioned above, the PMNO **1205** is configured to support multiple RAN vendors based on preferences of the managed service provider (MSP). When a customer is onboarding, in some embodiments, an MSP can select one or more RAN vendor templates based on deployment criteria (or criterion). Based on the selected RAN vendor template(s), the PMNO **1205** associates the RAN OAM server's IP address with the customer and site, according to some embodiments. Once the associations have been made, any subsequent changes to the 5G core and/or RAN deployment from a user portal would trigger the PMNO **1205**, in some embodiments, to invoke vendor-specific APIs against the OAM server endpoints **1210a-1214b**.

(72) The parameters assigned by the centralized management system include TACs for different customer premise locations, and DNNs for the different data networks operating in the SD-PMN. FIG. **13** conceptually illustrates an architecture diagram of a centrally managed SD-PMN of some embodiments after TACs and DNNs have been assigned. As shown, the diagram **1300** includes multiple sites **1330**, **1332**, and **1334** each having a respective UPF **1340**, **1342**, and **1344** deployed for the site and at least one respective access point **1350**, **1352**, and **1354**. In addition to the multiple sites, multiple multi-tenant 5G control planes **1320**, **1322**, and **1324** are deployed to PoPs **1360**, **1362**, and **1364** throughout the SD-PMN and that connect to a single core **1310** that is centrally managed by a PMN orchestrator **1305**. The single core **1310** is a master source for the control plane instances **1320-1324** and syncs with each control plane instance **1320-1324** as indicated. The diagram **1300** will be further described below by reference to FIG. **14**, which conceptually illustrates a process of some embodiments for centrally managing a SD-PMN. The process **1400** is performed in some embodiments by a centralized management server for the SD-PMN, such as the PMN orchestrator **1305** in the diagram **1300**.

(73) The process **1400** starts when for each physical location spanned by the SD-PMN, the process receives (at **1410**) a TAC defined for the physical location. In some embodiments, the PMN orchestrator **1305** is a server that provides a user interface (UI) through which a user (e.g., network administrator) can provide input such as TACs defined for physical locations spanned by the SD-PMN. The UI, in some embodiments, includes multiple selectable UI items for providing the input and configuring other aspects of the SD-PMN through, e.g., drop down menus, radio buttons, selection boxes, text fields, etc. For instance, in some embodiments, the UI includes a particular text field or set of text fields for defining TACs for each of the physical locations. The UI, in some embodiments, requires the TACs to follow a particular format, such as by using hexadecimal values that is two octets in length.

(74) In the diagram **1300**, each of the sites **1330-1334** includes a unique TAC. For example, the first site **1330** is assigned the TAC 315010:10008, the second site **1332** is assigned the TAC 315010:10009, and the last site **1334** is assigned the TAC 315010:10nnn. Each TAC is an identifier of the physical location area within the SD-PMN and is unique across all of the physical locations. The TACs, in some embodiments, are also associated with the access points **1350-1354** deployed in the respective physical locations such that the TACs can be used to identify a physical location and/or one or more access points in a physical location.

(75) For each data network running with the SD-PMN, the process **1400** receives (at **1420**) a DNN defined for the data network. In some embodiments, multiple VLANs (virtual local area networks) are implemented within the SD-PMN and are each assigned a respective DNN that spans all of the physical locations spanned by the SD-PMN. The UI provided by the management server of some embodiments includes multiple UI items for defining the DNNs in addition to the multiple UI items for defining the TACs as mentioned above. In some embodiments, a portion of the DNN is determined by the type of data network being named (e.g., “VLAN”), and a text field is provided to enable the user to further define the DNN by, e.g., adding a number or series of numbers.

(76) Because the data networks span all of the physical locations, each physical location of some embodiments may be associated with a particular data network and DNN or set of data networks and DNNs. For example, in the diagram **1300**, the first site **1330** is associated with DNNs VLAN2 and VLAN2, the second site **1332** is associated with DNNs VLAN 1 and VLAN3, and the last site **1334** is associated with DNNs VLAN1 and VLANn, as illustrated. Because each data network (and associated DNN) spans all of the physical locations, the VLAN1 is included in the DNNs associated with each site. While each site **1330-1334** is shown as having two associated DNNs, other sites in other embodiments can have more associated DNNs or fewer associated DNNs than illustrated.

(77) The process **1400** stores (at **1430**) the received TACs and DNNs in a core of the SD-PMN. The PMN orchestrator **1305** of some embodiments stores the received TACs and DNNs in the core **1310**, which syncs with the control plane instances **1320-1324** deployed to the PoPs **1360, 1362, and 1364**. Because each control plane instance **1320-1324** is the same as each other control plane instance **1320-1324**, each control plane instance **1320-1324** is aware of the TAC defined for and assigned to each physical site, as well as the DNNs defined for and assigned to each data network in the SD-PMN.

(78) The process **1400** assigns (at **1440**) TACs and DNNs to user devices as the user devices join the network. In some embodiments, the UPF assigns the new user device's traffic to a particular DNN or set of DNNs. When a user device moves to a new primary location, in some embodiments, an updated TAC or set of TACs is subsequently assigned to the user device based on the new location of the user device. For instance, a user device of some embodiments can join the SD-PMN while being primarily located at the first site **1330** in the diagram **1300**, and later change its primary location to the last site **1334**. As such, the TAC associated with the user device would be updated from 315010:10008 to 315010:10nnn in some such embodiments. Following **1440**, the process **1400** ends.

(79) In some embodiments, an SD-PMN may be implemented as a fixed wireless network for a particular geographic location. For example, a rural area might require internet access, while managed service providers (e.g., internet service providers, mobile network service providers, etc.) decline to deploy service for that rural area due to factors such as cost to the service provider. FIG. 15 conceptually illustrates a process of some embodiments for implementing an SD-PMN as a fixed wireless network for a particular geographic location. The process 1500 will be described below with references to FIG. 16, which conceptually illustrates the architecture of a fixed wireless network of some embodiments.

(80) The process 1500 starts when for each physical location in a set of physical locations within the particular geographic area, the process deploys (at 1510) an SD-WAN customer premise appliance. In the architecture diagram 1600, SD-WAN edge routers 1610 are deployed to homes 1615 within a particular geographic area for which the SD-PMN is being implemented. The SD-WAN edge routers 1610 are enabled with LTE (long term evolution), 4G, or 5G SIM, according to some embodiments. Each home premise within the geographic area for which the SD-PMN is being implemented, in some embodiments, will have an SD-WAN edge router to stretch the SD-PMN to that home premise.

(81) The process 1500 establishes (at 1520), via a RAN, a set of connections between each SD-WAN customer premise appliance and at least one physical access point deployed to the particular geographic area. Each of the SD-WAN edge routers 1610 at the home premises 1615 has a connection to the network of towers 1620 in the diagram 1600, as shown. The connections between the SD-WAN edge routers 1610 and the towers (i.e., access points) 1620 are LTE, 4G, or 5G connections via CBRS (citizens broadband radio service), in some embodiments.

(82) In some embodiments, the SD-WAN edge routers deployed to home premises (or other location types in the geographic area) are provided by the service provider of the SD-PMN and act as general residential broadband customer premise equipment (CPEs), while the access points (i.e. base stations, antennas, towers, etc.) deployed to the geographic area are provided by a third-party network equipment vendor and paid for by the entity for which the SD-PMN is being implemented. Examples of such entities can include corporations, educational campuses, and municipalities (e.g., towns, cities, etc.), according to some embodiments.

(83) The process 1500 connects (at 1530) each physical access point deployed to the particular geographic area to a central aggregation point to enable internet access for the particular geographic area. The network of towers 1620, for example, connects to the UPF 1630 that is part of the edge compute stack 1625. The access points in the network of towers 1620, in some embodiments, aggregate layer 2 (L2) traffic that terminates at the distributed UPF 1630. From the UPF 1630, the traffic goes through an aggregator SD-WAN edge router 1635 toward the SD-WAN gateway 1650 at the SASE PoP 1640, which also includes a 5G core 1655. The SASE PoP 1640 provides optimized internet connectivity, according to some embodiments.

(84) The process 1500 uses (at 1540) the established set of connections to provide SD-PMN service to the particular geographic area. That is, once the connections have been established, user devices at the home premises 1615 can access the internet 1660 through the series of connections that start from the SD-WAN edge routers 1610. The SD-PMN is controlled and managed as described in the embodiments above. Following 1540, the process 1500 ends.

(85) The fixed wireless solution described above differs from the architectures described by other embodiments of the invention in that the fixed wireless network 1600 uses an SD-WAN edge router as customer premise equipment to connect to the network, thus creating two layers of SD-WAN. The first layer provides the connection (e.g., a VCMP (VeloCloud multipath) tunnel) between the SD-WAN edge routers 1610 at the home premises 1615 (or other premises types for other entities) and the SD-WAN gateway 1650 at the SD-WAN/SASE PoP 1640 of the SD-PMN provider, and the second layer provides the connection between the SD-WAN edge router 1635 that is part of the edge compute stack 1625 and the SD-WAN gateway 1650 at the SD-WAN/SASE PoP 1640 of the

SD-PMN provider. Despite the two layers of SD-WAN, the SD-WAN edge router **1635** that is part of the edge compute stack **1625** does not utilize double tunneling and instead uses the first tunnel established by the SD-WAN edge routers **1610** on the underlay, while still protecting against failures on multiple WAN links and providing optimized fixed wireless access for WAN users. (86) The tunnel established from the SD-WAN edge routers **1610** to the SD-WAN gateway **1650** is optimized, in some embodiments, using DMPO, like in the other embodiments described above. In some embodiments, the SD-WAN edge router **1635** only sends traffic for which DMPO or any other optimization is desired to the SD-WAN gateway **1650**, and sends any other traffic to its destination without going through the SD-WAN gateway **1650**, while in other embodiments, all traffic is sent to its destination through the SD-WAN gateway **1650**. In some embodiments, sending all traffic through the SD-WAN gateway **1650** is desirable based on the 5G core **1655** for the SD-PMN also being located in the PoP **1640**.

(87) Some embodiments of the invention implement joint orchestration across an SD-PMN, SD-WAN, and edge compute stacks to enable customers to describe desired edge applications to be deployed alongside connectivity and QoS requirements, and to use the provided descriptions to orchestrate the edge application, connectivity, and QoS requirements across the SD-PMN, SD-WAN, and edge compute stack to yield the desired end-to-end connectivity and QoS for the desired edge application and any devices accessing the desired edge application. FIG. 17 conceptually illustrates a diagram of a joint orchestration platform of some embodiments that orchestrates applications end-to-end across an SD-PMN, SD-WAN, and edge compute stack.

(88) As shown, the joint orchestration platform diagram **1700** includes a PMN orchestrator (PMNO) **1705**, an edge compute stack (ECS) management system **1710**, a RAN/Core management system **1720**, an SD-WAN management plane **1730**, and a SASE management plane **1740**. Examples of an ECS management system, in some embodiments, include VMware Telco Cloud Automation (TCA) and VMware Tanzu Kubernetes Grid (TKG). An example of both the SD-WAN management plane and the SASE management plane are orchestrators, in some embodiments, is the VeloCloud Orchestrator (VCO).

(89) As shown, the PMNO **1705** has northbound intent-based APIs **1750** (application programming interfaces) to collect edge application connectivity requirements. The northbound APIs **1750** can include edge application connectivity requirements such as edge application workload compute, storage, and networking requirements; device groups that need connectivity to the application and at what QoS level; and any QoS requirements needed between the edge application and the cloud. The northbound intent-based APIs **1750** are defined by a user (e.g., network administrator) that manages the joint orchestration platform **1700**, according to some embodiments.

(90) After receiving these requirements, the PMNO **1705** takes these requirements and uses southbound APIs in some embodiments to deploy the workload on the ECS, make a subscriber group for devices that need connectivity to the edge application and configure appropriate data networking for that subscriber group (e.g., VLAN, QoS, etc.), and program business policies in the orchestrator (e.g., management server) for the SD-WAN. As illustrated, the intent-based API **1752** is sent to the ECS management system **1710**, the intent-based API **1754** is sent to the RAN/Core management system **1720**, the intent-based API **1756** is sent to the SD-WAN management plane **1730**, and the intent-based API **1758** is sent to the SASE management plane **1740**. As a result of the joint orchestration platform **1700**, end-to-end SLAs (service-level agreements) can be met for the entire system.

(91) FIG. 18 illustrates a joint orchestration platform diagram that includes a set of example intent-based APIs for orchestrating a video analytics application, in some embodiments, across an SD-PMN, SD-WAN, and edge compute stack. The joint orchestration platform **1800** includes a PMNO **1805**, ECS management system **1810**, RAN/Core management system **1820**, SD-WAN management plane **1830**, and SASE management plane **1840**. In this example, the PMNO **1805** receives (i.e., from a network administrator) an intent-based API **1850** that is defined for a video

analytics application “X”. The intent-based API **1850** includes a URL for the application's container image (i.e., an unchangeable, stand-alone, static file that includes executable code and well-defined assumptions about the application's run-time environment), specifications indicating the application requires high priority compute resources and high QoS connectivity between cameras and the application, specifications indicating the application requires high priority internet communication, and an indication that internet traffic associated with the application requires URL filtering.

(92) After the PMNO **1805** receives the intent-based API **1850**, the PMNO **1805** generates intent-based APIs for each of the management systems and planes **1810-1840** to implement and orchestrate the video analytics application “X”. Each requirement defined in the intent-based API **1850** corresponds to a different orchestration platform within the joint orchestration platform **1800**. For example, definitions relating to the application itself (e.g., the container image URL) and its compute resources are directed to the ECS management system, definitions regarding QoS requirements are directed to the RAN/Core management system, definitions regarding traffic priority are directed to the SD-WAN management plane, and definitions regarding URL filtering are directed to the SASE management plane, according to some embodiments.

(93) As such, based on the application container image URL and high priority compute resources defined in the intent-based API **1850**, the PMNO **1805** generates the intent-based API **1852** to direct the ECS management system **1810** to deploy the application image with required resources for the application. Based on the high QoS connectivity requirement between the cameras and the video analytics application defined in the intent-based API **1850**, the PMNO **1805** generates the intent-based API **1854** to direct the RAN/Core management system **1820** to define cameras as a subscriber group and apply the required QoS policy for the RAN and core. Based on the requirement for high priority Internet communication defined for the application in the intent-based API **1850**, the PMNO **1805** generates the intent-based API **1854** to direct the SD-WAN management plane **1830** to identify application traffic and add a business policy for prioritizing that traffic. Lastly, based on the URL filtering for Internet traffic defined in the intent-based API **1850**, the PMNO **1805** generates the intent-based API **1858** to direct the SASE management plane **1840** to create a URL filtering security policy for the application and service chain cloud web security (CWS). After each management system and plane **1810-1840** receives its respective intent-based API **1852-1858**, the video analytics application “X” is implemented across the SD-PMN.

(94) Many of the above-described features and applications are implemented as software processes that are specified as a set of instructions recorded on a computer-readable storage medium (also referred to as computer-readable medium). When these instructions are executed by one or more processing unit(s) (e.g., one or more processors, cores of processors, or other processing units), they cause the processing unit(s) to perform the actions indicated in the instructions. Examples of computer-readable media include, but are not limited to, CD-ROMs, flash drives, RAM chips, hard drives, EPROMs, etc. The computer-readable media does not include carrier waves and electronic signals passing wirelessly or over wired connections.

(95) In this specification, the term “software” is meant to include firmware residing in read-only memory or applications stored in magnetic storage, which can be read into memory for processing by a processor. Also, in some embodiments, multiple software inventions can be implemented as sub-parts of a larger program while remaining distinct software inventions. In some embodiments, multiple software inventions can also be implemented as separate programs. Finally, any combination of separate programs that together implement a software invention described here is within the scope of the invention. In some embodiments, the software programs, when installed to operate on one or more electronic systems, define one or more specific machine implementations that execute and perform the operations of the software programs.

(96) FIG. **19** conceptually illustrates a computer system **1900** with which some embodiments of the invention are implemented. The computer system **1900** can be used to implement any of the above-

described hosts, controllers, gateway, and edge forwarding elements. As such, it can be used to execute any of the above described processes. This computer system **1900** includes various types of non-transitory machine-readable media and interfaces for various other types of machine-readable media. Computer system **1900** includes a bus **1905**, processing unit(s) **1910**, a system memory **1925**, a read-only memory **1930**, a permanent storage device **1935**, input devices **1940**, and output devices **1945**.

(97) The bus **1905** collectively represents all system, peripheral, and chipset buses that communicatively connect the numerous internal devices of the computer system **1900**. For instance, the bus **1905** communicatively connects the processing unit(s) **1910** with the read-only memory **1930**, the system memory **1925**, and the permanent storage device **1935**.

(98) From these various memory units, the processing unit(s) **1910** retrieve instructions to execute and data to process in order to execute the processes of the invention. The processing unit(s) **1910** may be a single processor or a multi-core processor in different embodiments. The read-only-memory (ROM) **1930** stores static data and instructions that are needed by the processing unit(s) **1910** and other modules of the computer system **1900**. The permanent storage device **1935**, on the other hand, is a read-and-write memory device. This device **1935** is a non-volatile memory unit that stores instructions and data even when the computer system **1900** is off. Some embodiments of the invention use a mass-storage device (such as a magnetic or optical disk and its corresponding disk drive) as the permanent storage device **1935**.

(99) Other embodiments use a removable storage device (such as a floppy disk, flash drive, etc.) as the permanent storage device. Like the permanent storage device **1935**, the system memory **1925** is a read-and-write memory device. However, unlike storage device **1935**, the system memory **1925** is a volatile read-and-write memory, such as random access memory. The system memory **1925** stores some of the instructions and data that the processor needs at runtime. In some embodiments, the invention's processes are stored in the system memory **1925**, the permanent storage device **1935**, and/or the read-only memory **1930**. From these various memory units, the processing unit(s) **1910** retrieve instructions to execute and data to process in order to execute the processes of some embodiments.

(100) The bus **1905** also connects to the input and output devices **1940** and **1945**. The input devices **1940** enable the user to communicate information and select commands to the computer system **1900**. The input devices **1940** include alphanumeric keyboards and pointing devices (also called "cursor control devices"). The output devices **1945** display images generated by the computer system **1900**. The output devices **1945** include printers and display devices, such as cathode ray tubes (CRT) or liquid crystal displays (LCD). Some embodiments include devices such as touchscreens that function as both input and output devices **1940** and **1945**.

(101) Finally, as shown in FIG. **19**, bus **1905** also couples computer system **1900** to a network **1965** through a network adapter (not shown). In this manner, the computer **1900** can be a part of a network of computers (such as a local area network ("LAN"), a wide area network ("WAN"), or an Intranet), or a network of networks (such as the Internet). Any or all components of computer system **1900** may be used in conjunction with the invention.

(102) Some embodiments include electronic components, such as microprocessors, storage and memory that store computer program instructions in a machine-readable or computer-readable medium (alternatively referred to as computer-readable storage media, machine-readable media, or machine-readable storage media). Some examples of such computer-readable media include RAM, ROM, read-only compact discs (CD-ROM), recordable compact discs (CD-R), rewritable compact discs (CD-RW), read-only digital versatile discs (e.g., DVD-ROM, dual-layer DVD-ROM), a variety of recordable/rewritable DVDs (e.g., DVD-RAM, DVD-RW, DVD+RW, etc.), flash memory (e.g., SD cards, mini-SD cards, micro-SD cards, etc.), magnetic and/or solid state hard drives, read-only and recordable Blu-Ray® discs, ultra-density optical discs, any other optical or magnetic media, and floppy disks. The computer-readable media may store a computer program

that is executable by at least one processing unit and includes sets of instructions for performing various operations. Examples of computer programs or computer code include machine code, such as is produced by a compiler, and files including higher-level code that are executed by a computer, an electronic component, or a microprocessor using an interpreter.

(103) While the above discussion primarily refers to microprocessor or multi-core processors that execute software, some embodiments are performed by one or more integrated circuits, such as application-specific integrated circuits (ASICs) or field-programmable gate arrays (FPGAs). In some embodiments, such integrated circuits execute instructions that are stored on the circuit itself.

(104) As used in this specification, the terms “computer”, “server”, “processor”, and “memory” all refer to electronic or other technological devices. These terms exclude people or groups of people. For the purposes of the specification, the terms “display” or “displaying” mean displaying on an electronic device. As used in this specification, the terms “computer-readable medium,” “computer-readable media,” and “machine-readable medium” are entirely restricted to tangible, physical objects that store information in a form that is readable by a computer. These terms exclude any wireless signals, wired download signals, and any other ephemeral or transitory signals.

(105) While the invention has been described with reference to numerous specific details, one of ordinary skill in the art will recognize that the invention can be embodied in other specific forms without departing from the spirit of the invention. Thus, one of ordinary skill in the art would understand that the invention is not to be limited by the foregoing illustrative details, but rather is to be defined by the appended claims.

Claims

1. A method of providing access to a plurality of cloud-delivered services for a multi-tenant SD-PMN (software-defined private mobile network), the SD-PMN spanning at least a first site belonging to a first entity and a PMN-provider second site, the first entity's first site comprising a first portion of the SD-PMN that includes one or more physical access points for connecting a plurality of user devices at the first entity's first site to the SD-PMN and one or more data plane components of the SD-PMN, the method comprising: at a gateway router deployed in the PMN-provider second site: receiving a data message from a source located at the first entity's first site; determining that the received data message should be processed by a service chain that performs a set of one or more cloud-delivered services before the data message is forwarded to a destination of the data message; forwarding the data message to the service chain that performs the set of one or more cloud-delivered services for processing; and upon receiving the processed data message from the service chain, forwarding the processed data message to the destination of the data message.
2. The method of claim 1, wherein: the PMN-provider second site comprises a SASE (secure access service edge) PoP (point of presence); the service chain comprises a SASE service chain located in the SASE PoP; and forwarding the data message to the service chain comprises forwarding the data message to the SASE service chain in the SASE PoP.
3. The method of claim 1, wherein forwarding the data message to the service chain comprises forwarding the data message to a third site that hosts the service chain.
4. The method of claim 1, wherein the source of the data message comprises a user device from the plurality of user devices at the first entity's first site.
5. The method of claim 4, wherein receiving the data message from the source located at the first entity's first site comprises receiving the data message from an edge router deployed to the first entity's first site to perform data message forwarding within the first entity's first site and between the first entity's first site and a plurality of destinations external to the first entity's first site.
6. The method of claim 5, wherein the data message traverses a set of elements between the user device and the edge router deployed to the first entity's first set, the set of elements comprising at least one physical access point deployed at the first entity's first site and a user plane function

(UPF) deployed at the first entity's first site.

7. The method of claim 1, wherein the SD-PMN is a multi-tenant SD-PMN, wherein the first entity is one of a plurality of entities serviced by the SD-PMN.

8. The method of claim 7, wherein the first entity's first site is one of a plurality of sites belonging to the first entity.

9. The method of claim 8, wherein each entity in the plurality of entities serviced by the SD-PMN has at least one site.

10. The method of claim 1, wherein the set of one or more cloud-delivered services comprises a firewall, a secure web gateway, a zero-trust network access service, and a threat detection service.

11. A non-transitory machine readable medium storing a program for execution by a set of processing units, the program for a gateway router that provides access to a plurality of cloud-delivered services for a multi-tenant SD-PMN (software-defined private mobile network), the SD-PMN spanning at least a first site belonging to a first entity and a PMN-provider second site to which the gateway router is deployed, the first entity's first site comprising a first portion of the SD-PMN that includes one or more physical access points for connecting a plurality of user devices at the first entity's first site to the SD-PMN and one or more data plane components of the SD-PMN, the program comprising sets of instructions for: receiving a data message from a source located at the first entity's first site; determining that the received data message should be processed by a service chain that performs a set of one or more cloud-delivered services before the data message is forwarded to a destination of the data message; forwarding the data message to the service chain that performs the set of one or more cloud-delivered services for processing; and upon receiving the processed data message from the service chain, forwarding the processed data message to the destination of the data message.

12. The non-transitory machine readable medium of claim 11, wherein: the PMN-provider second site comprises a SASE (secure access service edge) PoP (point of presence); the service chain comprises a SASE service chain located in the SASE PoP; and the set of instructions forwarding the data message to the service chain comprises a set of instructions forwarding the data message to the SASE service chain in the SASE PoP.

13. The non-transitory machine readable medium of claim 11, wherein the set of instructions forwarding the data message to the service chain comprises a set of instructions forwarding the data message to a third site that hosts the service chain.

14. The non-transitory machine readable medium of claim 11, wherein the source of the data message comprises a user device from the plurality of user devices at the first entity's first site.

15. The non-transitory machine readable medium of claim 14, wherein the set of instructions for receiving the data message from the source located at the first entity's first site comprises a set of instructions for receiving the data message from an edge router deployed to the first entity's first site to perform data message forwarding within the first entity's first site and between the first entity's first site and a plurality of destinations external to the first entity's first site.

16. The non-transitory machine readable medium of claim 15, wherein the data message traverses a set of elements between the user device and the edge router deployed to the first entity's first set, the set of elements comprising at least one physical access point deployed at the first entity's first site and a user plane function (UPF) deployed at the first entity's first site.

17. The non-transitory machine readable medium of claim 11, wherein the SD-PMN is a multi-tenant SD-PMN, wherein the first entity is one of a plurality of entities serviced by the SD-PMN.

18. The non-transitory machine readable medium of claim 17, wherein the first entity's first site is one of a plurality of sites belonging to the first entity.

19. The non-transitory machine readable medium of claim 18, wherein each entity in the plurality of entities serviced by the SD-PMN has at least one site.

20. The non-transitory machine readable medium of claim 11, wherein the set of one or more

cloud-delivered services comprises a firewall, a secure web gateway, a zero-trust network access service, and a threat detection service.
