

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12399560
Kind Code	B2
Date of Patent	August 26, 2025
Inventor(s)	Deleeuw; William C.

Natural human-computer interaction for virtual personal assistant systems

Abstract

Technologies for natural language interactions with virtual personal assistant systems include a computing device configured to capture audio input, distort the audio input to produce a number of distorted audio variations, and perform speech recognition on the audio input and the distorted audio variants. The computing device selects a result from a large number of potential speech recognition results based on contextual information. The computing device may measure a user's engagement level by using an eye tracking sensor to determine whether the user is visually focused on an avatar rendered by the virtual personal assistant. The avatar may be rendered in a disengaged state, a ready state, or an engaged state based on the user engagement level. The avatar may be rendered as semitransparent in the disengaged state, and the transparency may be reduced in the ready state or the engaged state. Other embodiments are described and claimed.

Inventors:	Deleeuw; William C. (Beaverton, OR)
Applicant:	Intel Corporation (Santa Clara, CA)
Family ID:	1000008782202
Assignee:	INTEL CORPORATION (Santa Clara, CA)
Appl. No.:	18/793290
Filed:	August 02, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240411365 A1	Dec. 12, 2024

Related U.S. Application Data

continuation parent-doc US 18186680 20230320 US 12099651 child-doc US 18793290
continuation parent-doc US 17532816 20211122 US 11609631 20230321 child-doc US 18186680
continuation parent-doc US 16902967 20200616 US 11181980 20211123 child-doc US 17532816
continuation parent-doc US 16257347 20190125 US 10684683 20200616 child-doc US 16902967
continuation parent-doc US 15452131 20170307 US 10198069 20190205 child-doc US 16257347
division parent-doc US 14129435 US 9607612 20170328 WO PCT/US2013/041866 20130520
child-doc US 15452131

Publication Classification

Int. Cl.: **G06F3/01** (20060101); **G06T13/80** (20110101); **G06V10/143** (20220101); **G06V40/19** (20220101); **G10L15/02** (20060101); **G10L15/22** (20060101); **G10L15/30** (20130101); **G10L21/003** (20130101)

U.S. Cl.:

CPC **G06F3/013** (20130101); **G06T13/80** (20130101); **G06V10/143** (20220101); **G06V40/19** (20220101); **G10L15/02** (20130101); **G10L15/22** (20130101); G10L2015/025 (20130101); G10L2015/227 (20130101); G10L2015/228 (20130101); G10L15/30 (20130101); G10L21/003 (20130101)

Field of Classification Search

CPC: G06F (3/013); G06F (3/167); G06T (13/80); G06V (10/143); G06V (40/19); G10L (15/02); G10L (15/22); G10L (15/30); G10L (21/003); G10L (2015/025); G10L (2015/227); G10L (2015/228)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5799276	12/1997	Komissarchik	N/A	N/A
5809471	12/1997	Brodsky	N/A	N/A
6317711	12/2000	Muroi	N/A	N/A
6360202	12/2001	Bhadkamkar	N/A	N/A
6793673	12/2003	Kowalsky	N/A	N/A
6824562	12/2003	Mathis	N/A	N/A
6908478	12/2004	Alferness	N/A	N/A
6949122	12/2004	Adams	N/A	N/A
6960229	12/2004	Mathis	N/A	N/A
6964683	12/2004	Kowalsky	N/A	N/A
6976995	12/2004	Mathis	N/A	N/A
7179282	12/2006	Alferness	N/A	N/A
7309354	12/2006	Mathis	N/A	N/A
7311729	12/2006	Mathis	N/A	N/A
7316708	12/2007	Gordon	N/A	N/A
7351260	12/2007	Nieminen	N/A	N/A
7503931	12/2008	Kowalsky	N/A	N/A

7608102	12/2008	Adams	N/A	N/A
7635387	12/2008	Reuter	N/A	N/A
7674287	12/2009	Alferness	N/A	N/A
7828841	12/2009	Mathis	N/A	N/A
7828842	12/2009	Nieminen	N/A	N/A
7837729	12/2009	Gordon	N/A	N/A
7857846	12/2009	Alferness	N/A	N/A
7966177	12/2010	Geiger	N/A	N/A
8062358	12/2010	Mathis	N/A	N/A
8075608	12/2010	Gordon	N/A	N/A
8172898	12/2011	Alferness	N/A	N/A
8182529	12/2011	Gordon	N/A	N/A
8223156	12/2011	Boss	N/A	N/A
8439971	12/2012	Reuter	N/A	N/A
8990079	12/2014	Newman	N/A	N/A
9607612	12/2016	Deleeuw	N/A	N/A
10022011	12/2017	Norton	N/A	N/A
10062004	12/2017	Sugahara	N/A	N/A
10198068	12/2018	Holz	N/A	N/A
10198069	12/2018	Deleeuw	N/A	N/A
10684683	12/2019	Deleeuw	N/A	N/A
11181980	12/2020	Deleeuw	N/A	N/A
11609631	12/2022	Deleeuw	N/A	N/A
12099651	12/2023	Deleeuw	N/A	N/A
2003/0083538	12/2002	Adams	N/A	N/A
2003/0105520	12/2002	Alferness	N/A	N/A
2003/0144697	12/2002	Mathis	N/A	N/A
2003/0149566	12/2002	Levin	N/A	N/A
2003/0212453	12/2002	Mathis	N/A	N/A
2003/0225454	12/2002	Mathis	N/A	N/A
2003/0236569	12/2002	Mathis	N/A	N/A
2003/0236661	12/2002	Burges	N/A	N/A
2004/0010305	12/2003	Alferness	N/A	N/A
2004/0111095	12/2003	Gordon	N/A	N/A
2004/0127980	12/2003	Kowalsky	N/A	N/A
2004/0193260	12/2003	Alferness	N/A	N/A
2004/0243228	12/2003	Kowalsky	N/A	N/A
2004/0249452	12/2003	Adams	N/A	N/A
2005/0004799	12/2004	Lyudovyyk	N/A	N/A
2005/0021121	12/2004	Reuter	N/A	N/A
2005/0033582	12/2004	Gadd	N/A	N/A
2005/0096666	12/2004	Gordon	N/A	N/A
2005/0119673	12/2004	Gordon	N/A	N/A
2005/0125229	12/2004	Kurzweil	N/A	N/A
2005/0149180	12/2004	Mathis	N/A	N/A
2005/0187619	12/2004	Mathis	N/A	N/A
2005/0209690	12/2004	Mathis	N/A	N/A
2005/0216077	12/2004	Mathis	N/A	N/A
2005/0272969	12/2004	Alferness	N/A	N/A
2006/0020335	12/2005	Kowalsky	N/A	N/A

2006/0142854	12/2005	Alferness	N/A	N/A
2006/0145944	12/2005	Tarlton	N/A	N/A
2006/0167544	12/2005	Nieminen	N/A	N/A
2006/0173536	12/2005	Mathis	N/A	N/A
2007/0038451	12/2006	Cogne	N/A	N/A
2007/0055293	12/2006	Alferness	N/A	N/A
2007/0066879	12/2006	Mathis	N/A	N/A
2007/0225972	12/2006	Kim	N/A	N/A
2008/0096533	12/2007	Manfredi	N/A	N/A
2008/0097594	12/2007	Mathis	N/A	N/A
2008/0109059	12/2007	Gordon	N/A	N/A
2008/0140191	12/2007	Mathis	N/A	N/A
2008/0319542	12/2007	Nieminen	N/A	N/A
2009/0150156	12/2008	Kennewick	N/A	N/A
2009/0271189	12/2008	Agapi	N/A	N/A
2010/0079467	12/2009	Boss	N/A	N/A
2010/0100175	12/2009	Reuter	N/A	N/A
2010/0156781	12/2009	Fahn	N/A	N/A
2010/0168847	12/2009	Alferness	N/A	N/A
2011/0015928	12/2010	Odell	N/A	N/A
2011/0035000	12/2010	Nieminen	N/A	N/A
2011/0066234	12/2010	Gordon	N/A	N/A
2011/0175810	12/2010	Markovic	N/A	N/A
2012/0154277	12/2011	Bar-Zeev	N/A	N/A
2012/0257035	12/2011	Larsen	N/A	N/A
2012/0295708	12/2011	Hernandez-Abrego	N/A	N/A
2013/0005443	12/2012	Kosta	N/A	N/A
2013/0005458	12/2012	Kosta	N/A	N/A
2013/0005482	12/2012	Kosta	N/A	N/A
2013/0006709	12/2012	Kosta	N/A	N/A
2014/0184550	12/2013	Hennessey	N/A	N/A
2014/0225901	12/2013	Shuster	N/A	N/A
2014/0267311	12/2013	Evertt	N/A	N/A
2016/0063989	12/2015	Deleeuw	N/A	N/A
2016/0210503	12/2015	Yin	N/A	N/A
2016/0260436	12/2015	Lemay	N/A	N/A
2017/0039751	12/2016	Tong	N/A	N/A
2017/0177080	12/2016	Deleeuw	N/A	N/A
2019/0187787	12/2018	White	N/A	N/A
2019/0227623	12/2018	Deleeuw	N/A	N/A
2019/0369748	12/2018	Hindi	N/A	N/A
2021/0072825	12/2020	Deleeuw	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
1134074	12/1995	CN	N/A
1164801	12/1996	CN	N/A
1343968	12/2001	CN	N/A
101292282	12/2007	CN	N/A
101939740	12/2010	CN	N/A

101952818	12/2010	CN	N/A
102129292	12/2010	CN	N/A
102591016	12/2011	CN	N/A
102749991	12/2011	CN	N/A
102814814	12/2011	CN	N/A
2012221498	12/2011	JP	N/A

OTHER PUBLICATIONS

“Intelligent personal assistant,” Wikipedia, The Free Encyclopedia, retrieved from:

<http://en.wikipedia.org/w/index.php?title=Intelligent.sub.-personal.-sub.--assistant&oldid=520490140>, edited Oct. 29, 2012, 2 pages. cited by applicant

“Automated online assistant,” Wikipedia, The Free Encyclopedia, retrieved from

<http://en.wikipedia.org/w/index.php?title=Automated.sub.-onlin-e.sub.-assistant&oldid=521739395>, edited Nov. 6, 2012, 3 pages. cited by applicant

“Eye tracking,” Wikipedia, The Free Encyclopedia, retrieved from:

<http://en.wikipedia.org/w/index.php?title=Eye.sub-tracking&oldid=52-3123367>, edited Nov. 15, 2012, 7 pages. cited by applicant

“Speech recognition,” Wikipedia, The Free Encyclopedia, retrieved from:

<http://en.wikipedia.org/w/index.php?title=Speech.sub.-recognition&oldid=528092164>, edited Dec. 15, 2012, 12 pages. cited by applicant

International Searching Authority, “International Search Report and Written Opinion,” issued in connection with International Patent Application No. PCT/US2013/041866, dated Feb. 28, 2014, 19 pages. cited by applicant

The International Bureau, “International Preliminary Report on Patentability,” issued in connection with International Patent Application No. PCT/US2013/041866, dated Nov. 24, 2015, 12 pages. cited by applicant

United States Patent and Trademark Office, “Non-Final Office Action,” issued in connection with U.S. Appl. No. 14/129,435, dated Feb. 16, 2016, 9 pages. cited by applicant

United States Patent and Trademark Office, “Final Office Action,” issued in connection with U.S. Appl. No. 14/129,435, dated Jun. 23, 2016, 16 pages. cited by applicant

United States Patent and Trademark Office, “Notice of Allowance and Fee(s) Due,” issued in connection with U.S. Appl. No. 14/129,435, dated Nov. 16, 2016, 13 pages. cited by applicant

Chinese Patent Office, English translation of “Search Report,” issued in connection with Chinese Patent Application No. 2013800757640 dated Mar. 14, 2018, 2 pages. cited by applicant

Chinese Patent Office, “First Office Action,” issued in connection with Chinese Patent Application No. 201380075764, dated Mar. 22, 2018, 24 pages (includes English translation). cited by applicant

United States Patent and Trademark Office, “Non-Final Office Action,” issued in connection with U.S. Appl. No. 15/452,131, dated May 17, 2018, 6 pages. cited by applicant

Chinese Patent Office, English translation of “Supplemental Search Report,” issued in connection with Chinese Patent Application No. 2013800757640 dated Aug. 8, 2018, 2 pages (includes English translation). cited by applicant

Chinese Patent Office, English translation of “Second Office Action,” issued in connection with Chinese Patent Application No. 201380075764, dated Sep. 5, 2018, 5 pages. cited by applicant

United States Patent and Trademark Office, “Notice of Allowance and Fee(s) Due,” issued in connection with U.S. Appl. No. 15/452,131, dated Sep. 25, 2018, 7 pages. cited by applicant

United States Patent and Trademark Office, “Notice of Allowance and Fee(s) Due,” issued in connection with U.S. Appl. No. 15/452,131, dated Oct. 17, 2018, 2 pages. cited by applicant

Chinese Patent Office, “Notification to Grant Patent Right for Invention,” issued in connection with Chinese Patent Application No. 2013800757640, dated Mar. 25, 2019, 4 pages (includes English translation). cited by applicant

United States Patent and Trademark Office, “Non-Final Office Action,” issued in connection with U.S. Appl. No. 16/257,347, dated May 2, 2019, 6 pages. cited by applicant

United States Patent and Trademark Office, “Final Office Action,” issued in connection with U.S. Appl. No. 16/257,347, dated Oct. 29, 2019, 7 pages. cited by applicant

United States Patent and Trademark Office, “Notice of Allowance and Fee(s) Due,” issued in connection with U.S. Appl. No. 16/257,347, dated Feb. 12, 2020, 7 pages. cited by applicant

United States Patent and Trademark Office, “Corrected Notice of Allowability,” issued in connection with U.S. Appl. No. 16/257,347, dated Apr. 9, 2020, 2 pages. cited by applicant

United States Patent and Trademark Office, “Non-Final Office Action,” issued in connection with U.S. Appl. No. 16/902,967, dated Mar. 22, 2021, 10 pages. cited by applicant

United States Patent and Trademark Office, “Notice of Allowance and Fee(s) due,” issued in connection with U.S. Appl. No. 16/902,967, dated Jul. 12, 2021, 7 pages. cited by applicant

United States Patent and Trademark Office, “Notice of Allowance,” issued in U.S. Appl. No. 16/902,967 on Oct. 20, 2021, 2 pages. cited by applicant

United States Patent and Trademark Office, “Non-Final Action” issued in U.S. Appl. No. 17/532,816 on Jun. 17, 2022 (8 pages). cited by applicant

National Intellectual Property Administration, P.R. China, “Notification to Grant Patent Right for Invention,” issued in connection with Chinese Patent Application No. 201811329915.1, dated Sep. 8, 2022, 7 pages (English translation included). cited by applicant

United States Patent and Trademark Office, “Notice of Allowance” issued in U.S. Appl. No. 17/532,816 on Nov. 16, 2022 (8 pages). cited by applicant

United States Patent and Trademark Office, “Non-Final Office Action,” issued in connection with U.S. Appl. No. 18/186,680, dated Feb. 1, 2024, 8 pages. cited by applicant

United States Patent and Trademark Office, “Notice of Allowance and Fee(s) Due,” issued in connection with U.S. Appl. No. 18/186,680, dated May 24, 2024, 7 pages. cited by applicant

Primary Examiner: Sarma; Abhishek

Attorney, Agent or Firm: HANLEY, FLIGHT & ZIMMERMAN, LLC

Background/Summary

CROSS REFERENCE TO RELATED APPLICATIONS (1) This patent arises from a continuation of U.S. patent application Ser. No. 18/186,680, entitled “NATURAL HUMAN-COMPUTER INTERACTION FOR VIRTUAL PERSONAL ASSISTANT SYSTEMS,” and filed on Mar. 20, 2023. U.S. patent application Ser. No. 18/186,680 is a continuation of U.S. patent application Ser. No. 17/532,816, now U.S. Pat. No. 11,609,631, entitled “NATURAL HUMAN-COMPUTER INTERACTION FOR VIRTUAL PERSONAL ASSISTANT SYSTEMS,” and filed on Nov. 22, 2021. U.S. patent application Ser. No. 17/532,816 is a continuation of U.S. patent application Ser. No. 16/902,967, now U.S. Pat. No. 11,181,980, entitled “NATURAL HUMAN-COMPUTER INTERACTION FOR VIRTUAL PERSONAL ASSISTANT SYSTEMS,” and filed on Jun. 16, 2020. U.S. patent application Ser. No. 16/902,967 is a continuation of U.S. patent application Ser. No. 16/257,347, now U.S. Pat. No. 10,684,683, entitled “NATURAL HUMAN-COMPUTER INTERACTION FOR VIRTUAL PERSONAL ASSISTANT SYSTEMS,” and filed on Jan. 25, 2019. U.S. patent application Ser. No. 16/257,347 is a continuation of U.S. patent application Ser. No. 15/452,131, now U.S. Pat. No. 10,198,069, entitled “NATURAL HUMAN-COMPUTER INTERACTION FOR VIRTUAL PERSONAL ASSISTANT SYSTEMS,” and filed on Mar. 7, 2017. U.S. patent application Ser. No. 15/452,131 is a divisional of U.S. patent application Ser. No.

14/129,435, now U.S. Pat. No. 9,607,612, entitled "NATURAL HUMAN-COMPUTER INTERACTION FOR VIRTUAL PERSONAL ASSISTANT SYSTEMS," and filed on Dec. 26, 2013. U.S. patent application Ser. No. 14/129,435 is a national stage entry under 35 USC § 371 (b) of International Patent Application No. PCT/US2013/041866, entitled "NATURAL HUMAN-COMPUTER INTERACTION FOR VIRTUAL PERSONAL ASSISTANT SYSTEMS," and filed on May 20, 2013. U.S. patent application Ser. No. 18/186,680, U.S. patent application Ser. No. 17/532,816; U.S. patent application Ser. No. 16/902,967; U.S. patent application Ser. No. 16/257,347; U.S. patent application Ser. No. 15/452,131; U.S. patent application Ser. No. 14/129,435; and International Patent Application No. PCT/US2013/041866 are hereby incorporated by reference in their entireties. Priority to each of U.S. patent application Ser. No. 18/186,680, U.S. patent application Ser. No. 17/532,816; U.S. patent application Ser. No. 16/902,967; U.S. patent application Ser. No. 16/257,347; U.S. patent application Ser. No. 15/452,131; U.S. patent application Ser. No. 14/129,435; and International Patent Application No. PCT/US2013/041866 is hereby claimed.

BACKGROUND

(1) Virtual personal assistants are artificial intelligence systems that perform tasks on a computing device in response to natural-language requests from a user. For example, a virtual personal assistant may handle calendaring, reminders, and messaging tasks for the user. To interact with the virtual personal assistant, the user typically enters a pre-defined input sequence on the computing device, for example pressing a dedicated hardware button or speaking a predefined code word. The user may enter natural-language requests through conventional text input or through speech recognition.

(2) To further facilitate natural interaction, many virtual personal assistants display a humanlike character, also known as an avatar, to serve as a main point of interaction with the user. The avatar may occupy or obscure a significant portion of the display of the computing device. Further, the avatar may interfere with use of other applications on the computing device, particularly when the user did not intend to activate the avatar. Even when displaying a humanlike avatar, typical systems may not fully model natural human interaction, and instead may require conventional human-computer interactions such as button presses, mouse clicks, or the like.

(3) Speech recognition systems convert spoken utterances of the user into computer-readable representations of text. Typical speech recognition systems attempt to determine a single most-likely speech recognition result for a given audio input. Such systems may filter out noise or otherwise attempt to enhance the audio input signal in order to improve speech recognition results. Some systems may provide a small number of alternative results; however, these results are typically only slight variations on each other. Typical speech recognition engines may be implemented as components of a local computing device, or as services provided by a server computing device.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The concepts described herein are illustrated by way of example and not by way of limitation in the accompanying figures. For simplicity and clarity of illustration, elements illustrated in the figures are not necessarily drawn to scale. Where considered appropriate, reference labels have been repeated among the figures to indicate corresponding or analogous elements.

(2) FIG. 1 is a simplified block diagram of at least one embodiment of a computing device for natural interaction with a virtual personal assistant;

(3) FIG. 2 is a simplified block diagram of at least one embodiment of an environment of the computing device of FIG. 1;

- (4) FIG. 3 is a simplified block diagram of at least one embodiment of system for natural interaction with a virtual personal assistant;
- (5) FIG. 4 is a simplified flow diagram of at least one embodiment of a method for introducing audio distortion to improve speech recognition that may be executed by the computing device of FIGS. 1 and 2;
- (6) FIG. 5 is a simplified flow diagram of at least one embodiment of a method for introducing audio distortion to improve speech recognition that may be executed by the speech recognition server of FIG. 4; and
- (7) FIGS. 6A and 6B are a simplified flow diagram of at least one embodiment of a method for natural interaction with a virtual personal assistant that may be executed by the computing device of FIGS. 1 and 2.

DETAILED DESCRIPTION OF THE DRAWINGS

- (8) While the concepts of the present disclosure are susceptible to various modifications and alternative forms, specific embodiments thereof have been shown by way of example in the drawings and will be described herein in detail. It should be understood, however, that there is no intent to limit the concepts of the present disclosure to the particular forms disclosed, but on the contrary, the intention is to cover all modifications, equivalents, and alternatives consistent with the present disclosure and the appended claims.
- (9) References in the specification to “one embodiment,” “an embodiment,” “an illustrative embodiment,” etc., indicate that the embodiment described may include a particular feature, structure, or characteristic, but every embodiment may or may not necessarily include that particular feature, structure, or characteristic. Moreover, such phrases are not necessarily referring to the same embodiment. Further, when a particular feature, structure, or characteristic is described in connection with an embodiment, it is submitted that it is within the knowledge of one skilled in the art to effect such feature, structure, or characteristic in connection with other embodiments whether or not explicitly described.
- (10) The disclosed embodiments may be implemented, in some cases, in hardware, firmware, software, or any combination thereof. The disclosed embodiments may also be implemented as instructions carried by or stored on a transitory or non-transitory machine-readable (e.g., computer-readable) storage medium, which may be read and executed by one or more processors. A machine-readable storage medium may be embodied as any storage device, mechanism, or other physical structure for storing or transmitting information in a form readable by a machine (e.g., a volatile or non-volatile memory, a media disc, or other media device).
- (11) In the drawings, some structural or method features may be shown in specific arrangements and/or orderings. However, it should be appreciated that such specific arrangements and/or orderings may not be required. Rather, in some embodiments, such features may be arranged in a different manner and/or order than shown in the illustrative figures. Additionally, the inclusion of a structural or method feature in a particular figure is not meant to imply that such feature is required in all embodiments and, in some embodiments, may not be included or may be combined with other features.
- (12) Referring now to FIG. 1, an illustrative computing device **100** for natural interaction with a virtual personal assistant includes a processor **120**, an I/O subsystem **122**, and memory **124**. The computing device **100** captures audio input and obtains speech recognition results from a speech recognition engine that are based on distortions applied to the audio input. The applied distortions may result in multiple, semantically distinct variations of the audio input. Supplying multiple variations may allow for the speech recognition engine to produce many more potential speech recognition results, which in turn may increase speech recognition accuracy. The speech recognition engine may operate on the computing device **100**, or in some embodiments, on a remote speech recognition server as discussed in more detail below. The speech recognition results may be used to control a virtual personal assistant. The virtual personal assistant models the

engagement level of the user by tracking the user's visual focus and/or by interpreting the user's speech. Modeling user engagement may allow the virtual personal assistant to engage in more-natural human interactions; for example, the virtual personal assistant may better determine when it is being addressed by the user or when the user has moved on to some other task. The virtual personal assistant may represent the engagement level of the user by adjusting the size, position, and/or transparency of an avatar on a display screen. Rendering the engagement level in this manner may also facilitate natural interaction by allowing the user to better understand the state of the virtual personal assistant without unnecessarily interrupting the user's other work.

(13) The computing device **100** may be embodied as any type of device capable of performing the functions described herein. For example, the computing device **100** may be embodied as, without limitation, a smartphone, a cellular phone, a tablet computer, a notebook computer, a laptop computer, a desktop computer, a distributed computing system, a multiprocessor system, a consumer electronic device, a smart appliance, and/or any other computing device capable of recognizing spoken user commands. As shown in FIG. **1**, the illustrative computing device **100** includes the processor **120**, the I/O subsystem **122**, the memory **124**, and a data storage device **126**. Of course, the computing device **100** may include other or additional components, such as those commonly found in a portable computer (e.g., various input/output devices), in other embodiments. Additionally, in some embodiments, one or more of the illustrative components may be incorporated in, or otherwise from a portion of, another component. For example, the memory **124**, or portions thereof, may be incorporated in the processor **120** in some embodiments.

(14) The processor **120** may be embodied as any type of processor currently known or developed in the future and capable of performing the functions described herein. For example, the processor may be embodied as a single or multi-core processor(s), digital signal processor, microcontroller, or other processor or processing/controlling circuit. Similarly, the memory **124** may be embodied as any type of volatile or non-volatile memory or data storage currently known or developed in the future and capable of performing the functions described herein. In operation, the memory **124** may store various data and software used during operation of the computing device **100** such as operating systems, applications, programs, libraries, and drivers. The memory **124** is communicatively coupled to the processor **120** via the I/O subsystem **122**, which may be embodied as circuitry and/or components to facilitate input/output operations with the processor **120**, the memory **124**, and other components of the computing device **100**. For example, the I/O subsystem **122** may be embodied as, or otherwise include, memory controller hubs, input/output control hubs, firmware devices, communication links (i.e., point-to-point links, bus links, wires, cables, light guides, printed circuit board traces, etc.) and/or other components and subsystems to facilitate the input/output operations. In some embodiments, the I/O subsystem **122** may form a portion of a system-on-a-chip (SoC) and be incorporated, along with the processor **120**, the memory **124**, and other components of the computing device **100**, on a single integrated circuit chip.

(15) The data storage **126** may be embodied as any type of device or devices configured for short-term or long-term storage of data such as, for example, memory devices and circuits, memory cards, hard disk drives, solid-state drives, or other data storage devices. The data storage **126** may store program and data files relating to the virtual personal assistant, and may serve as temporary or permanent storage for audio data captured by the computing device **100**.

(16) The computing device **100** further includes a display **128**, an audio sensor **130**, and an eye tracking sensor **132**. The display **128** of the computing device **100** may be embodied as any type of display capable of displaying digital information such as a liquid crystal display (LCD), a light emitting diode (LED), a plasma display, a cathode ray tube (CRT), or other type of display device. In some embodiments, the display **128** may be coupled to a touch screen to receive user input.

(17) The audio sensor **130** may be embodied as any sensor capable of capturing audio signals such as a microphone, a line input jack, an analog-to-digital converter (ADC), or other type of audio sensor. The audio sensor **130** may be used by the computing device **100** to detect speech commands

uttered by the user, as described below.

(18) The eye tracking sensor **132** may be embodied as any one or more sensors capable of determining an area on the display **128** of the computing device **100** on which the user's eyes are focused. For example, the eye tracking sensor **132** may be embodied as a digital camera or a digital depth camera capable of tracking the focus of the user's gaze. In other embodiments, the eye tracking sensor **132** may be embodied as active infrared emitters and infrared detectors capable of tracking the user's eye movements over time. In those embodiments, the eye tracking sensor **132** may capture the infrared light reflected off of various internal and external features of the user's eye and thereby calculate the direction of the user's gaze. The eye tracking sensor **132** may also be capable of determining the position of the user's head in three-dimensional space. In some embodiments, an eye tracking sensor **132** such as a depth camera may be capable of determining head position data directly. In other embodiments, the eye tracking sensor **132** may be used with another sensor such as a video camera to calculate the position of the user's head.

(19) In some embodiments, the computing device **100** may also include one or more peripheral devices **134**. The peripheral devices **134** may include any number of additional sensors, input/output devices, interface devices, and/or other peripheral devices. For example, in some embodiments, the peripheral devices **134** may include a touch screen, graphics circuitry, keyboard, mouse, speaker system, interface devices, and/or other input/output devices. In some embodiments, the peripheral devices **134** may be used along with the eye tracking sensor **132** and/or the audio sensor **130** to determine the user's engagement level. As another example, in some embodiments, the peripheral devices **134** may include a communication circuit, device, or collection thereof capable of enabling communications between the computing device **100** and other remote servers and/or devices.

(20) Referring now to FIG. 2, in one embodiment, the computing device **100** establishes an environment **200** during operation. The illustrative embodiment **200** includes a speech recognition module **202**, an audio input module **204**, a virtual personal assistant **208**, and an engagement module **214**. The various modules of the environment **200** may be embodied as hardware, firmware, software, or a combination thereof.

(21) The speech recognition module **202** is configured to perform speech recognition on audio input data received from the audio input module **204**. The speech recognition module **202** ranks and filters speech recognition results to produce a single result or a ranked list of likely results. The speech recognition module **202** may use a speech recognition grammar provided by an application such as the virtual personal assistant **208** to rank and filter speech recognition results. In some embodiments, the speech recognition module **202** may recognize speech in a dictation or free speech mode. The dictation or free speech mode may use a full natural language vocabulary and grammar to recognize results, and thus may produce additional likely speech recognition results.

(22) The audio input module **204** captures audio input data from the audio sensor **130** and applies audio distortions to the audio input data to produce multiple variations of the audio input. The audio distortions may modify amplitude, frequency, duration, and/or other characteristics of the audio input to produce semantic variation among the distorted audio variations. The audio input module **204** provides the distorted audio variations to the speech recognition module **202**. In some embodiments, those functions may be performed by sub-modules, for example, by a distortion module **206**. Additionally, in some embodiments, the functionality of the speech recognition module **202** and/or the distortion module **206** may be performed by a remote server, for example by a cloud service, as described below in connection with FIG. 3.

(23) The virtual personal assistant **208** responds to spoken user commands and displays an avatar on the display **128** to provide information on the status of the virtual personal assistant **208**. The virtual personal assistant **208** may maintain a speech recognition grammar defining spoken commands that may be accepted from the user, including command vocabulary and syntax. The avatar is a character or other visual representation of the virtual personal assistant **208**. The avatar

may include human-like characteristics such as facial features or a human form. Those human-like features may facilitate natural interaction with the user. In some embodiments, those functions may be performed by sub-modules, for example a command module **210** or an avatar module **212**. In some embodiments, the virtual personal assistant **208** may be used without speech recognition; that is, the virtual personal assistant **208** may respond to non-speech input such as typed input or input gestures.

(24) The engagement module **214** determines the user's level of engagement with the virtual personal assistant **208** based on sensor data received from the eye tracking sensor **132** and/or the audio sensor **130**. For example, the engagement module **214** may determine the level of engagement based on how long or how often the user's eyes focus on the avatar. In some embodiments, the engagement module **214** may also analyze speech recognition results from the speech recognition module **202** to determine the user's level of engagement. The engagement module **214** provides the engagement level to the virtual personal assistant **208**, allowing the virtual personal assistant **208** to modify the avatar accordingly.

(25) Referring now to FIG. 3, in some embodiments, the speech recognition module **202** and/or the distortion module **206** may be embodied in a remote speech recognition server **300**. The speech recognition server **300** is configured to provide services including performing speech recognition analysis on audio input transmitted from the computing device **100** over a network **302**. The speech recognition server **300** may be embodied as any type of server computing device, or collection of devices, capable of performing the functions described herein. As such, the speech recognition server **300** may include components and features similar to the computing device **100** such as a processor, I/O subsystem, memory, data storage, communication circuitry, and various peripheral devices, which are not illustrated in FIG. 3 for clarity of the present description. Further, the speech recognition server **300** may be embodied as a single server computing device or a collection of servers and associated devices. For example, in some embodiments, the speech recognition server **300** may be embodied as a “virtual server” formed from multiple computing devices distributed across a network and operating in a public or private cloud. Accordingly, although the speech recognition server **300** is illustrated in FIG. 3 and described below as embodied as a single server computing device, it should be appreciated that the speech recognition server **300** may be embodied as multiple devices cooperating together to facilitate the functionality described below.

(26) As discussed above, the speech recognition module **202** and, in some embodiments, the audio distortion module **206** are established on the speech recognition server **300** (i.e., rather than, or in addition to, the computing device **100**). The computing device **100** includes the audio input module **204**, the virtual personal assistant **208**, the engagement module **214**, and in some embodiments, the distortion module **206**. The various modules of the speech recognition server **300** and the computing device **100** perform the same functions as the modules described above in connection with FIG. 2, and may be embodied as hardware, firmware, software, or a combination thereof.

(27) The illustrative computing device **100** of FIG. 3 further includes a speech analysis module **304**. The speech analysis module **304** is configured to transmit the audio input, and in some embodiments, the distorted audio variations to the speech recognition server **300**. The speech analysis module **304** is also configured to receive speech recognition results from the speech recognition server **300** and may select a result from the speech recognition results based on contextual information.

(28) In embodiments in which the speech recognition module **202** and the distortion module **206** are located on the speech recognition server **300**, the audio input module **204** of the computing device **100** captures audio input from the audio sensor **130**, and the speech analysis module **304** sends the audio input to the speech recognition server **300**. In such embodiments, the distortion module **206** of the speech recognition server **300** receives the audio input and applies audio distortions to the audio input data to produce multiple variations of the audio input, as described above with regard to the computing device **100**. The audio input and the distorted audio variations

are provided to the speech recognition module **202** of the speech recognition server **300**, which performs speech recognition as described above. The speech recognition module **202** subsequently transmits the speech recognition results back to the computing device **100**. The speech recognition module **202** may transmit a list of speech recognition results, or may transmit a single result.

(29) Alternatively, in embodiments in which only the speech recognition module **202** is located on the speech recognition server **300**, the audio input module **204** of the computing device **100** captures audio input from the audio sensor **130** and the distortion module **206** of the computing device **100** applies audio distortions to the audio input data to produce multiple variations of the audio input, as described above. In such embodiments, the computing device **100** transmits the audio input and the distorted audio variations to the speech recognition server **300**. The audio input and the distorted audio variations are received by the speech recognition module **202** of the speech recognition server **300**, which performs speech recognition as described above. The speech recognition module **202** transmits the speech recognition results back to the computing device **100**. The speech recognition module **202** may transmit a list of speech recognition results, or may transmit a single result.

(30) Referring now to FIG. **4**, in use, the computing device **100** may execute a method **400** for introducing audio distortion to improve speech recognition. The method **400** begins with block **402**, in which the computing device **100** captures audio input using the audio sensor **130**. The audio input may be stored in any format useable for further analysis and manipulation, including compressed or uncompressed formats. Audio capture may be initiated based on a request from the user or based on a request from a component of the computing device **100**. For example, audio capture may be initiated when the user activates or engages the virtual personal assistant **208**, as further described below.

(31) In block **404**, in some embodiments the computing device **100** may pre-process the captured audio input. As part of pre-processing, the computing device **100** may perform speech recognition on the captured audio input. The speech recognition results from pre-processing may be used to control distortion of the audio input, as described below.

(32) In block **406**, the computing device **100** distorts the audio input to produce multiple distorted variations of the audio input. The audio input is distorted in ways that may produce different and semantically varied versions of the audio input. The distortions may modify, for example, the amplitude, timing, pitch, or any other salient characteristic of the audio input. The computing device **100** may apply the distortions individually or in various combinations. Each variation may be stored as separate audio data, or the computing device **100** may apply distortions to the original audio input as needed.

(33) For example, in some embodiments, the computing device **100** may remove quiet audio segments from the audio input in block **408**. To do so, the computing device **100** may identify internal segments of the audio input (that is, segments that are not at the beginning or the end of the audio input) having an amplitude below a threshold amplitude and delete those internal segments from the audio input. After removing a quiet segment, segments of the audio input that were formerly on either side of the quiet segment are appended together or otherwise run together. Therefore, this distortion may eliminate pauses between utterances spoken by the user, which may combine the utterances into a single word.

(34) In some embodiments, the computing device **100** may expand quiet audio segments in the audio input in block **410**. To do so, the computing device **100** may identify segments of the audio input having an amplitude below a threshold amplitude and increase the duration of those segments. After expanding the quiet segments, neighboring segments of the audio input are spaced further apart. Therefore, this distortion may increase the duration of pauses between utterances spoken by the user, which may split the utterances into two distinct words.

(35) In some embodiments, the computing device **100** may insert pauses at one or more phonetic split points in the audio input identified in pre-processing in block **412**. As with quiet audio

segment expansion, this distortion may split utterances of the user into distinct words. Unlike quiet audio expansion, this distortion does not require an existing pause or quiet audio segment.

(36) In some embodiments, in block **414**, the computing device **100** may modify the playback speed of the audio input; that is, the computing device **100** may speed up or slow down the audio input. In some embodiments, in block **416**, the computing device **100** may apply other audio transformations to the audio input. For example, the computing device **100** may alter the pitch of the audio input or mix the audio input with random noise. Such distortions may produce variation in speech recognition results. Of course, the computing device **100** may apply additional or other distortions to the audio input.

(37) In block **418**, in embodiments in which the speech recognition module **202** is located on the remote speech recognition server **300**, the computing device **100** may transmit the audio input and the multiple distorted audio variations to the speech recognition server **302**. As described above, the speech recognition server may provide speech recognition services. Offloading speech recognition tasks to the speech recognition server **302** may improve efficiency for a computing device **100** having limited computational power, for example, a smartphone.

(38) After applying the various distortions to produce multiple variations of the audio input, the computing device **100** obtains speech recognition results based on the audio input and all of the multiple distorted audio variations in block **420**. For example, in embodiments in which the speech recognition module **202** is located on the mobile computing device **100**, the device **100** may perform speech recognition locally. In those embodiments, the computing device **100** may re-use or reference speech recognition results created while pre-processing the audio input in block **404** and perform speech recognition in block **420** on only the distorted audio variations. In other embodiments in which the speech recognition module **202** is located on the speech recognition server **300**, the computing device **100** may receive speech recognition results from the speech recognition server **300**. However the speech recognition results are obtained, the speech recognition process produces multiple speech recognition results; that is, multiple potential interpretations of the captured audio input. In some embodiments, the speech recognition process may produce a large number of speech recognition results; for example, hundreds or thousands of speech recognition results may be produced. In some embodiments, the speech recognition results may be generated using a conventional speech recognition engine supplied with the multiple variations of the audio input repeatedly or in parallel.

(39) In block **422**, in some embodiments, the computing device **100** may determine semantically relevant speech recognition results from among the full speech recognition results produced in block **420**. In other words, the computing device **100** may analyze the speech recognition results to find results that it “understands” and to reject results that do not “make sense.” For example, an application of the computing device **100** such as the virtual personal assistant **208** may include a speech recognition grammar. The speech recognition grammar may include a set of actions, data objects, and other commands understood by the application. The computing device **100** may determine semantically relevant speech recognition results by accepting only those results that satisfy the speech recognition grammar.

(40) In block **424**, the computing device **100** selects a result from speech recognition results based on contextual information. In some embodiments, the result may be selected from the smaller set of semantically relevant speech recognition results determined in block **422**. The selected result may be the speech recognition result most relevant to the current context of the user and/or the computing device **100**. The most relevant speech recognition result is most likely to be the result intended by the user. The contextual information includes any information that may reveal the intent of the user, and may include the state of any applications currently executing on the computing device **100**, such as web browsers, productivity applications, or the virtual personal assistant **208**. The contextual information may also include contextual information associated with the user, such as a calendar, contact list, email account, or other personalized data. The contextual

information may further include basic contextual information of the computing device **100** such as date, time, or location. Although illustrated as selecting a single result from the speech recognition results, in some embodiments the computing device **100** may produce a list of speech recognition results, which list may be sorted based on the contextual information. After selecting a result, the method **400** is completed and control may be passed to a calling method that may operate on the speech recognition results. For example, as described below, the virtual personal assistant **208** may act on the speech recognition results.

(41) Referring now to FIG. 5, in embodiments in which the speech recognition module **202** and/or the distortion module **206** are located on the speech recognition server **300**, the speech recognition server **300** may execute a method **500** for introducing audio distortion to improve speech recognition. The method **500** begins with block **502**, in which the speech recognition server **302** determines whether a request for speech recognition has been received from a computing device **100**. If not, the method **500** loops back to block **502** to continue listening for speech recognition requests. However, if a request has been received, the method **500** advances to block **504** in which the speech recognition server **300** receives audio input from computing device **100**. The audio input was previously captured by the computing device **100** using the audio sensor **130**. The audio input may be stored in any format useable for further analysis and manipulation, including compressed or uncompressed formats.

(42) In block **506**, in some embodiments the speech recognition server **300** may pre-process the received audio input. As part of pre-processing, the speech recognition server **300** may perform speech recognition on the received audio input. The speech recognition results from pre-processing may be used to control distortion of the audio input, as described below.

(43) In block **508**, the speech recognition server **300** distorts the audio input to produce multiple distorted variations of the audio input. The audio input is distorted in ways that may produce different and semantically varied versions of the audio input. The distortions may modify, for example, the amplitude, timing, pitch, or any other salient characteristic of the audio input. The speech recognition server **300** may apply the distortions individually or in various combinations. Each variation may be stored as separate audio data, or the speech recognition server **300** may apply distortions to the original audio input as needed.

(44) For example, in some embodiments, the speech recognition server **300** may remove quiet audio segments from the audio input in block **510**, expand quiet audio segments in the audio input in block **512**, insert pauses at one or more phonetic split points in the audio input identified in pre-processing in block **514**, modify the playback speed of the audio input in block **516**, and/or apply other audio transforms to the audio input in block **518** as described in detail above with regard to computing device **100** and block **406** of method **400**. Of course, the speech recognition server **300** may apply additional or other distortions to the audio input. Regardless, after applying the various distortions to produce multiple variations of the audio input, the speech recognition server **300** performs speech recognition on the audio input and all of the multiple distorted audio variations in block **520**. In some embodiments, the speech recognition server **300** may re-use or reference speech recognition results created while pre-processing the audio input in block **506** and perform speech recognition in block **520** on only the distorted audio variations. The speech recognition process produces multiple speech recognition results; that is, multiple potential interpretations of the captured audio input. In some embodiments, the speech recognition process may produce a large number of speech recognition results; for example, hundreds or thousands of speech recognition results may be produced. In some embodiments, the speech recognition results may be generated using a conventional speech recognition engine supplied with the multiple variations of the audio input repeatedly or in parallel.

(45) In block **522**, in some embodiments, the speech recognition server **300** may determine semantically relevant speech recognition results from among the full speech recognition results produced in block **520**. In other words, the speech recognition server **300** may analyze the speech

recognition results to find results that it “understands” and to reject results that do not “make sense.” For example, an application of the computing device **100** such as the virtual personal assistant **208** may include a speech recognition grammar. The speech recognition grammar may include a set of actions, data objects, and other commands understood by the application. The speech recognition server **300** may determine semantically relevant speech recognition results by accepting only those results that satisfy the speech recognition grammar of the virtual personal assistant **208**.

(46) In block **524**, in some embodiments, the speech recognition server **300** may select a result from speech recognition results based on contextual information. For example, in some embodiments, the result may be selected from the smaller set of semantically relevant speech recognition results determined in block **522**. The selected result may be the speech recognition result most relevant to the current context of the user and/or the computing device **100**. The most relevant speech recognition result is most likely to be the result intended by the user. The contextual information may include any information that may reveal the intent of the user, and may include the state of any applications currently executing on the computing device **100**, such as web browsers, productivity applications, or the virtual personal assistant **208**. The contextual information may also include contextual information associated with the user, such as a calendar, contact list, email account, or other personalized data. The contextual information may further include basic contextual information of the computing device **100** such as date, time, or location. Although illustrated as selecting a single result from the speech recognition results, in some embodiments the speech recognition server **302** may produce a list of speech recognition results, which list may be sorted based on the contextual information.

(47) In block **526**, the speech recognition server **302** transmits the speech recognition result or speech recognition results back to the computing device **100**. The computing device **100** may then use the speech recognition result or speech recognition results to control an application such as the virtual personal assistant **208**. After transmitting, the method **500** loops back to block **502** to listen for further speech recognition requests from the computing device **100**.

(48) Referring now to FIG. **6A**, in use, the computing device **100** may execute a method **600** for natural interaction with the virtual personal assistant **208**. The method **600** begins with block **602**, in which an avatar of the virtual personal assistant **208** is displayed in a disengaged state on the display **128** of the computing device **100**. As described above, the avatar is a visual representation of the virtual personal assistant **208**. The avatar includes human-like or anthropomorphic features that may facilitate natural interaction with the user. For example, the avatar may include an animated human or human-like face including an eye or a pair of eyes. When in the disengaged state, the avatar is displayed in a manner that indicates that the virtual personal assistant **208** is idle and/or not actively listening to instructions. For example, the avatar may be represented as sleeping or looking away from the user.

(49) In some embodiments, in block **604**, the avatar may be displayed in the disengaged state as semitransparent, allowing background applications of the computing device **100** to shine through the avatar. When semitransparent, the avatar may remain visible to the user while still allowing all of the display **128** to be useable by other applications and at least partially visible to the user. The semitransparent appearance may be accomplished through alpha-blending the avatar and the other applications, compositing a scene, or by any other comparable technique for semitransparency. In some embodiments, in block **606**, the avatar may be displayed in the disengaged state in a relatively small size or at a position away from the currently active application of the computing device **100**. For example, the avatar may be rendered in a smaller size in a corner of the display **128**, allowing the user to continue working in the active application without distraction. In other embodiments, the computing device **100** may render the avatar in the disengaged state as completely invisible. In those embodiments, all of the display **128** may be available for other applications; however, the user may not be presented with any visual cues indicating the virtual

personal assistant **208** is available.

(50) In block **608**, the computing device **100** monitors for activation of the virtual personal assistant **208** by the user. In some embodiments, in block **610**, the computing device **100** may receive eye tracking data from the eye tracking sensor **132**. The computing device **100** interprets the eye tracking data to determine the position of the user's gaze on the display **128**. The user may indicate activation of the virtual personal assistant **208** by focusing on the position of the avatar currently displayed in the disengaged state. The computing device **100** may filter the eye tracking data to remove brief or spurious glances at the avatar. In some embodiments, in block **612**, the computing device **100** may receive audio input from the audio sensor **130**. The computing device **100** interprets the audio input to determine whether the user has uttered a code word for activating the virtual personal assistant **208**. For example, the code word may be embodied as the name of the virtual personal assistant **208**, or a common word such as “assistant” or “computer.”

(51) In block **614**, the computing device **100** determines whether the user has requested activation of the virtual personal assistant **208**. As described above, activation may be requested when the user's gaze has focused on the avatar for a length of time longer than a certain threshold, or when the code word has been detected. If the user has not requested activation, the method **600** loops back to block **602**, keeping the avatar in the disengaged state and continuing to monitor for activation. If the user has requested activation, the method **600** advances to block **616**.

(52) In block **616**, the computing device **100** displays the avatar in a ready state. The ready state indicates to the user that the virtual personal assistant **208** is available and ready to respond to user input. In block **618**, in some embodiments the computing device **100** may render the avatar as making eye contact with the user. Eye contact is a powerful cue that may naturally convey to the user that the virtual personal assistant **208** is ready for input. In some embodiments, the computing device **100** may simply render the avatar as facing straight out of the display **128**, because the user is typically positioned in front of the display **128**. In some embodiments, the computing device **100** may use eye tracking data received from the eye tracking sensor **132** to determine the location of the user's eyes and render the avatar as looking at the user's eyes. The computing device **100** may render eye contact using anthropomorphic cues to simulate natural interaction; for example, the computing device **100** may render the avatar as periodically glancing away from the user, which may make the user less likely to interpret the avatar as staring at the user. In some embodiments, in block **620** the avatar may be displayed in the ready state by decreasing the transparency of the avatar; that is, by making the avatar appear more solid. While still allowing background applications to shine through, the increasingly solid appearance of the avatar may indicate that the virtual personal assistant **208** is ready to receive commands. In block **622**, in some embodiments the computing device **100** may display the avatar in the ready state by adjusting the position, size, and/or visibility of the avatar. For example, the avatar may be moved toward the active application on the display **128**, may be increased in size, or may be made visible.

(53) While the avatar is displayed in the ready state, in block **624** the computing device **100** determines the user's engagement level, that is, the level of interest the user is exhibiting in the avatar. By determining the user's engagement level, the computing device **100** may determine in a natural manner whether or not the user intended to activate the virtual personal assistant **208**. In some embodiments, in block **626** the computing device **100** may receive eye tracking data from the eye tracking sensor **132**. The computing device **100** may analyze the eye tracking data to determine whether the user has visually focused on the avatar. Visually focusing on the avatar may indicate a relatively high engagement level, and focusing away from the avatar may indicate relatively a low engagement level. The computing device **100** may require the user to visually focus on the avatar for a period of time greater than a threshold time, in order to filter out spurious glances. Similarly, the computing device **100** may ignore short glances away from the avatar to improve recognition performance, essentially applying a low-pass filter to the eye tracking data. In some embodiments, in block **628** the computing device **100** may receive audio input from the audio sensor **130**. The

computing device **100** may analyze the audio input to determine whether the user is addressing the virtual personal assistant **208**. In some embodiments, the computing device **100** may perform speech recognition on the audio input.

(54) In block **630**, the computing device **100** determines whether the user is engaged with the avatar. The computing device **100** may determine whether user is engaged by comparing the user engagement level determined in block **624** with a threshold engagement level. If not the user is determined not to be engaged, the method **600** loops back to block **602** to render the avatar in the disengaged state. To summarize that interaction, the avatar has unobtrusively indicated to the user that it is ready for interaction, determined that the user is not engaged with the avatar, and then retreated away from the user's attention. Thus, that interaction may simulate a natural human interaction. Referring again to block **630**, if the user is engaged with the avatar, the method **600** advances to block **632**, illustrated in FIG. **6B**.

(55) Referring now to FIG. **6B**, in block **632** the computing device **100** displays the avatar in an engaged state. The engaged state indicates to the user that the virtual personal assistant **208** is actively interpreting commands issued by the user. For example, in some embodiments, the computing device **100** may render the avatar as making eye contact with the user in block **634**. As described above in connection with block **618**, the computing device **100** may receive eye tracking data from the eye tracking sensor **132** and render the avatar as looking toward the user's eyes. Additionally or alternatively, in some embodiments, the computing device **100** may decrease the transparency of the avatar in block **636**. For example, the computing device **100** may render the avatar as fully opaque in some embodiments. Further, in some embodiments, the computing device **100** may adjust the size and/or position of the avatar in block **638**. For example, the avatar may be rendered close to or in front of the currently-active application on the display **128**, or the avatar may be increased in size. In some embodiments, although the avatar is displayed as opaque and in front of other applications on the display **128**, the avatar does not intercept or otherwise interfere with user input, allowing the user to continue working with the computing device **100**.

(56) While displaying the avatar in the engaged state, execution of the method **600** proceeds in parallel to blocks **640** and **644**. In block **640**, the computing device **100** performs speech recognition on audio input received from the audio sensor **130** while the avatar is in the engaged state. In some embodiments, the computing device **100** may perform a more accurate or more computationally-intense speech recognition method while in the engaged state, because it is likely that the user is directly addressing the virtual personal assistant **208**. For example, the computing device **100** may perform the speech recognition method of introducing distortion described above in connection with FIG. **4**, or any other speech recognition technique.

(57) In block **642**, the computing device **100** determines whether or not a command has been received that may be executed by the virtual personal assistant **208**. The computing device **100** may apply the results of speech recognition determined in block **640** to a speech recognition grammar of the virtual personal assistant **208** to determine whether a valid command has been issued. If no command has been received, the method **600** loops back to block **640** to continue performing speech recognition. If a command has been received, the method **600** advances to block **652** as described below.

(58) As described above, the method **600** executes block **644** in parallel with block **640**. In block **644**, the computing device **100** monitors the user engagement level while in the engaged state. As described above in connection with block **624**, the computing device **100** determines the level of interest the user is exhibiting in the avatar, which may allow for more natural interactions. In some embodiments, in block **646**, the computing device **100** may receive eye tracking data from the eye tracking sensor **132**. As described above with respect to block **626**, the computing device **100** may determine the engagement level based on whether or not the user's eyes are focused on the avatar. In some embodiments, in block **648** the computing device **100** may receive audio input from the audio sensor **130**. As described above in connection with block **628**, the computing device **100** may

analyze the audio input to determine whether the user is addressing the virtual personal assistant **208**. In some embodiments, the computing device **100** may use speech recognition results from block **640** to determine whether the user is addressing the virtual personal assistant **208**.

(59) In block **650**, the computing device **100** determines whether the user is engaged with the avatar. As described above in connection with block **630**, the computing device **100** may compare the user engagement level determined in block **644** with a threshold engagement level. If the user is engaged with the avatar, the method **600** loops back to block **644** to continue monitoring the user engagement level. To summarize that interaction, if the user remains actively engaged with the avatar, for example, by engaging in eye contact with the avatar, the avatar also remains in the engaged state. That interaction may simulate a natural interaction as when holding a conversation. If in block **650** the computing device **100** determines that the user is not engaged, the method **600** loops back to block **602** of FIG. **6A** to render the avatar in the disengaged state. To summarize that interaction, if after some time of engagement the user is no longer engaged, for example by turning back to other work, then the avatar also leaves, or begins to leave, the engaged state. That interaction may simulate a natural interaction as when a person completes a conversation.

(60) When the method **600** advances from block **642** to block **652** or loops back from block **650** to block **602**, the computing device **100** encounters a task synchronization boundary, illustrated in FIG. **6B** by a thick line. The task synchronization boundary ensures that only one task of method **600** remains active in the computing device **100**. For example, when advancing from block **642** to block **652** based on a received speech command, the computing device **100** may terminate the task executing block **644**, causing the computing device **100** to stop monitoring the user engagement level. Similarly, when looping back from block **650** to block **602**, the computing device **100** may terminate the task executing block **640**, causing the computing device **100** to stop performing speech recognition. Further, although illustrated as being performed in parallel, in other embodiments the tasks of performing speech recognition and monitoring the user engagement level may be performed sequentially or in an interleaved manner.

(61) In block **652**, the computing device **100** displays the avatar in a working state. The working state indicates to the user that the virtual personal assistant **208** is currently executing a task. In some embodiments, the working state includes a representation of the task being performed, for example, an application icon or a representation of the avatar performing a task. In some embodiments, the avatar displayed in the working state may be similar or identical to the avatar displayed in the disengaged state; that is, the avatar displayed in the working state may be unobtrusive and may not interfere with the user performing other work on the computing device **100**. In some embodiments, in block **654** the computing device **100** may increase the transparency of the avatar on the display **128**. In some embodiments, in block **656**, the computing device **100** may adjust the size and/or position of the avatar. For example, the computing device **100** may decrease the size of the avatar or move the avatar away from a currently active application on the display **128**.

(62) In block **658**, while the avatar is displayed in the working state the computing device **100** executes the command received from the user. The command may be executed by the virtual personal assistant **208** or by various other applications of the computing device **100** controlled by the virtual personal assistant **208**. Further, in some embodiments, while the avatar is displayed in the working state, the computing device **100** may monitor for activation of the assistant by the user, similar to as described above in connection with FIG. **408**. Monitoring for activation may allow the user to interrupt the currently executing command or initiate a new command.

(63) After completion of the command, in block **660** the computing device **100** determines whether to notify the user of completion. For example, if execution of the command produced displayable results or error messages, the computing device **100** may determine to notify the user. If the computing device **100** determines to notify the user, the method **600** loops back to block **616** of FIG. **6A** to display the avatar in the ready state. To summarize that interaction, the avatar works on

a task unobtrusively in the background for some time and then notifies the user of completion in a natural manner. The user may then interact with the avatar as described above to receive the notification. If in block **660** the computing device **100** determines not to notify the user, the method **600** loops back to block **602** of FIG. **6A** to display the avatar in the disengaged state and await further activation. To summarize that interaction, the avatar works on a task unobtrusively in the background for some time and then returns to the disengaged state, indicating to the user in a natural manner that the virtual personal assistant **208** is free for further interactions.

(64) In the preceding illustrative embodiment, the user engagement has been described as a binary value-either engaged or not engaged. However, in other embodiments, the user engagement level may be measured on a continuum. In those embodiments, the avatar may be displayed with properties that reflect the value of the user engagement level. For example, the transparency of the avatar may be smoothly adjusted over a range of values to indicate the user engagement level. Similarly, the size or position of the avatar may also be smoothly adjusted over a range of values to indicate the user engagement level.

EXAMPLES

(65) Illustrative examples of the devices, systems, and methods disclosed herein are provided below. An embodiment of the devices, systems, and methods may include any one or more, and any combination of, the examples described below.

(66) Example 1 includes a computing device for speech recognition, the computing device comprising an audio sensor; an audio input module to capture audio input using the audio sensor; and distort the audio input to produce a plurality of distorted audio variations; and a speech recognition module to perform speech recognition on the audio input and each of the distorted audio variations to produce a plurality of speech recognition results; and select a result from the speech recognition results based on contextual information.

(67) Example 2 includes the subject matter of Example 1, and wherein to distort the audio input comprises to remove an internal segment of the audio input.

(68) Example 3 includes the subject matter of any of Examples 1 and 2, and wherein the internal segment of the audio input comprises a segment having an amplitude with a predefined relationship to an amplitude threshold.

(69) Example 4 includes the subject matter of any of Examples 1-3, and wherein to distort the audio input comprises to expand a length of a segment of the audio input having an amplitude with a predefined relationship to an amplitude threshold.

(70) Example 5 includes the subject matter of any of Examples 1-4, wherein the segment having an amplitude with the predefined relationship to the amplitude threshold comprises a segment having an amplitude below the amplitude threshold.

(71) Example 6 includes the subject matter of any of Examples 1-5, and wherein to distort the audio input comprises to insert a pause at a phonetic split point of the audio input identified by performing speech recognition on the audio input.

(72) Example 7 includes the subject matter of any of Examples 1-6, and wherein to distort the audio input comprises at least one of to: (i) speed up the audio input, (ii) slow down the audio input, (iii) adjust a pitch of the audio input, or (iv) introduce noise to the audio input.

(73) Example 8 includes the subject matter of any of Examples 1-7, and wherein the plurality of speech recognition results comprises at least one hundred speech recognition results.

(74) Example 9 includes the subject matter of any of Examples 1-8, and further comprising one or more applications having a speech recognition grammar; wherein the speech recognition module is further to determine semantically relevant results of the speech recognition results based on the speech recognition grammar of the one or more applications; and wherein to select the result from the speech recognition results comprises to select a result from the semantically relevant results.

(75) Example 10 includes the subject matter of any of Examples 1-9, and wherein the one or more applications comprise a virtual personal assistant.

(76) Example 11 includes a computing device for speech recognition, the computing device comprising: an audio sensor; an audio input module to: capture audio input using the audio sensor; and distort the audio input to produce a plurality of distorted audio variations; and a speech analysis module to: transmit the audio input and the distorted audio variations from the computing device to a speech recognition module; receive a plurality of speech recognition results from the speech recognition module based on the audio input and the distorted audio variations; and select a result from the speech recognition results based on contextual information.

(77) Example 12 includes the subject matter of Example 11, and wherein the speech recognition module is located on a server remote from the computing device.

(78) Example 13 includes the subject matter of any of Examples 11 and 12, and wherein to distort the audio input comprises to remove an internal segment of the audio input.

(79) Example 14 includes the subject matter of any of Examples 11-13, and wherein the internal segment of the audio input comprises a segment having an amplitude with a predefined relationship to an amplitude threshold.

(80) Example 15 includes the subject matter of any of Examples 11-14, and wherein to distort the audio input comprises to expand a length of a segment of the audio input having an amplitude with a predefined relationship to an amplitude threshold.

(81) Example 16 includes the subject matter of any of Examples 11-15, and wherein the segment having an amplitude with the predefined relationship to the amplitude threshold comprises a segment having an amplitude below the amplitude threshold.

(82) Example 17 includes the subject matter of any of Examples 11-16, and wherein to distort the audio input comprises to insert a pause at a phonetic split point of the audio input identified by performing speech recognition on the audio input.

(83) Example 18 includes the subject matter of any of Examples 11-17, and wherein to distort the audio input comprises at least one of to: (i) speed up the audio input, (ii) slow down the audio input, (iii) adjust a pitch of the audio input, or (iv) introduce noise to the audio input.

(84) Example 19 includes the subject matter of any of Examples 11-18, and wherein the plurality of speech recognition results comprises at least one hundred speech recognition results.

(85) Example 20 includes a speech recognition server for speech recognition, the speech recognition server comprising: a distortion module to: receive audio input captured by a computing device; and distort the audio input to produce a plurality of distorted audio variations; and a speech recognition module to: perform speech recognition on the audio input and each of the distorted audio variations to produce a plurality of speech recognition results; and transmit the plurality of speech recognition results to the computing device.

(86) Example 21 includes the subject matter of Example 20, and wherein to distort the audio input comprises to remove an internal segment of the audio input.

(87) Example 22 includes the subject matter of any of Examples 20 and 21, and wherein the internal segment of the audio input comprises a segment having an amplitude with a predefined relationship to an amplitude threshold.

(88) Example 23 includes the subject matter of any of Examples 20-22, and wherein to distort the audio input comprises to expand a length of a segment of the audio input having an amplitude with a predefined relationship to an amplitude threshold.

(89) Example 24 includes the subject matter of any of Examples 20-23, and wherein the segment having the amplitude with the predefined relationship to the amplitude threshold comprises a segment having an amplitude below the amplitude threshold.

(90) Example 25 includes the subject matter of any of Examples 20-24, and wherein to distort the audio input comprises to insert a pause at a phonetic split point of the audio input identified by performing speech recognition on the audio input.

(91) Example 26 includes the subject matter of any of Examples 20-25, and wherein to distort the audio input comprises one of to: (i) speed up the audio input, (ii) slow down the audio input, (iii)

adjust a pitch of the audio input, or (iv) introduce noise to the audio input.

(92) Example 27 includes the subject matter of any of Examples 20-26, and wherein the plurality of speech recognition results comprises at least one hundred speech recognition results.

(93) Example 28 includes a computing device comprising: a display; a virtual personal assistant to display an avatar of the virtual personal assistant on the display in a disengaged state, a ready state, and an engaged state; and an engagement module to: determine whether a user of the computing device has requested activation of the virtual personal assistant while the avatar is displayed in the disengaged state; and determine an engagement level of the user while the avatar is displayed in the ready state; wherein the virtual personal assistant is to: display the avatar in the ready state in response to a determination that the user has requested activation of the virtual personal assistant; display the avatar in the engaged state in response to a determination that the user has an engagement level greater than a threshold level; and display the avatar in the disengaged state in response to a determination that the user has an engagement level less than the threshold level.

(94) Example 29 includes the subject matter of Example 28, and wherein the engagement module is further to monitor the engagement level of the user while the avatar is displayed in the engaged state.

(95) Example 30 includes the subject matter of any of Examples 28-29, and further comprising an eye tracking sensor, wherein to determine whether the user has requested activation of the virtual personal assistant comprises to: receive eye tracking data from the eye tracking sensor; and determine whether the user has focused on the avatar based on the eye tracking data.

(96) Example 31 includes the subject matter of any of Examples 28-30, and further comprising an audio sensor, wherein to determine whether the user has requested activation of the virtual personal assistant comprises to: receive audio input from the audio sensor; and detect a code word uttered by the user based on the audio input.

(97) Example 32 includes the subject matter of any of Examples 28-31, and further comprising an eye tracking sensor, wherein to determine the engagement level of the user comprises to: receive eye tracking data from the eye tracking sensor; determine, based on the eye tracking data, whether the user has visually focused on the avatar for a period of time greater than a threshold time; determine that the user has an engagement level greater than the threshold level in response to a determination that the user has visually focused on the avatar for a period of time greater than the threshold time; and determine that the user has an engagement level less than the threshold level in response to a determination that the user has not visually focused on the avatar for a period of time greater than the threshold time.

(98) Example 33 includes the subject matter of any of Examples 28-32, and wherein to determine whether the user has visually focused on the avatar for a period of time greater than the threshold time comprises to ignore glances away from the avatar for a second period of time less than a second threshold time.

(99) Example 34 includes the subject matter of any of Examples 28-33, and further comprising an audio sensor, wherein to determine the engagement level of the user comprises to receive audio input from the audio sensor.

(100) Example 35 includes the subject matter of any of Examples 28-34, and wherein: to display the avatar in the ready state comprises to display an anthropomorphic representation of eye contact of the avatar with the user; and to display the avatar in the engaged state comprises to display an anthropomorphic representation of eye contact of the avatar with the user.

(101) Example 36 includes the subject matter of any of Examples 28-35, and further comprising an eye tracking sensor, wherein to display the anthropomorphic representation of eye contact comprises to: receive at least one of eye tracking data or head position data from the eye tracking sensor; and display an anthropomorphic representation of eyes of the avatar following the user.

(102) Example 37 includes the subject matter of any of Examples 28-36, and wherein: to display the avatar of the virtual personal assistant in the disengaged state comprises to display the avatar as

semitransparent, to allow a user interface element of the computing device to shine through the avatar; to display the avatar in the ready state comprises to decrease a transparency of the avatar; and to display the avatar in the engaged state comprises one of to: decrease the transparency of the avatar or eliminate the transparency of the avatar.

(103) Example 38 includes the subject matter of any of Examples 28-37, and wherein: to display the avatar of the virtual personal assistant in the disengaged state comprises to display the avatar at a position on the display away from an active application of the computing device; to display the avatar in the ready state comprises to move the avatar on the display to a position closer to an active application of the computing device; and to display the avatar in the engaged state comprises to move the avatar on the display to a position over an active application of the computing device without preventing input from the user to the active application.

(104) Example 39 includes the subject matter of any of Examples 28-38, and further comprising an eye tracking sensor, wherein: to display the avatar in the ready state comprises to move the avatar on the display to a position closer to a position where the user is focused, based on eye tracking data received from the eye tracking sensor; and to display the avatar in the engaged state comprises to move the avatar on the display to a position closer to a position where the user is focused, based on eye tracking data received from the eye tracking sensor.

(105) Example 40 includes the subject matter of any of Examples 28-39, and wherein: to display the avatar of the virtual personal assistant in the disengaged state comprises to display the avatar as invisible; and to display the avatar in the ready state comprises to display the avatar as visible.

(106) Example 41 includes the subject matter of any of Examples 28-40, and wherein the virtual personal assistant is further to: perform speech recognition while the avatar is in the engaged state; determine whether the user has issued a command based on the speech recognition; and display the avatar in a working state in response to a determination that the user has issued the command.

(107) Example 42 includes the subject matter of any of Examples 28-41, and wherein to display the avatar in the working state comprises at least one of to: (i) increase a transparency of the avatar or (ii) move the avatar on the display to a position away from an active application of the computing device.

(108) Example 43 includes a method for speech recognition on a computing device, the method comprising: capturing audio input using an audio sensor of the computing device; distorting the audio input to produce a plurality of distorted audio variations; performing speech recognition on the audio input and each of the distorted audio variations to produce a plurality of speech recognition results; and selecting a result from the speech recognition results based on contextual information.

(109) Example 44 includes the subject matter of Example 43, and wherein distorting the audio input comprises removing an internal segment of the audio input.

(110) Example 45 includes the subject matter of any of Examples 43 and 44, and wherein removing the internal segment of the audio input comprises removing a segment of the audio input having an amplitude with a predefined relationship to an amplitude threshold.

(111) Example 46 includes the subject matter of any of Examples 43-45, and wherein removing the internal segment comprises removing the segment having an amplitude below the amplitude threshold.

(112) Example 47 includes the subject matter of any of Examples 43-46, and wherein distorting the audio input comprises expanding a length of a segment of the audio input having an amplitude with a predefined relationship to an amplitude threshold.

(113) Example 48 includes the subject matter of any of Examples 43-47, and wherein expanding the length of the segment comprises expanding a length of the segment having amplitude below the amplitude threshold.

(114) Example 49 includes the subject matter of any of Examples 43-48, and wherein distorting the audio input comprises inserting a pause at a phonetic split point of the audio input identified by

performing speech recognition on the audio input.

(115) Example 50 includes the subject matter of any of Examples 43-49, and wherein distorting the audio input comprises performing at least one of: (i) speeding up the audio input, (ii) slowing down the audio input, (iii) adjusting a pitch of the audio input, or (iv) introducing noise to the audio input.

(116) Example 51 includes the subject matter of any of Examples 43-50, and wherein performing speech recognition on the audio input and the distorted audio variations to produce the plurality of speech recognition results comprises performing speech recognition on the audio input and the distorted audio variations to produce at least one hundred speech recognition results.

(117) Example 52 includes the subject matter of any of Examples 43-51, and further comprising determining semantically relevant results of the speech recognition results based on a speech recognition grammar of one or more applications of the computing device; wherein selecting the result from the speech recognition results comprises selecting a result from the semantically relevant results.

(118) Example 53 includes the subject matter of any of Examples 43-52, and wherein determining the semantically relevant results based on the speech recognition grammar of the one or more applications comprises determining the semantically relevant results based on a speech recognition grammar of a virtual personal assistant of the computing device.

(119) Example 54 includes a method for speech recognition on a computing device, the method comprising: capturing audio input using an audio sensor of the computing device; distorting, on the computing device, the audio input to produce a plurality of distorted audio variations; transmitting the audio input and the distorted audio variations from the computing device to a speech recognition module; receiving a plurality of speech recognition results from the speech recognition module based on the audio input and the distorted audio variations; and selecting, on the computing device, a result from the speech recognition results based on contextual information.

(120) Example 55 includes the subject matter of Example 54, and wherein: transmitting the audio input and the distorted audio variations to the speech recognition module comprises transmitting the audio input and the distorted audio variations to a speech recognition module located on a server remote from the computing device; and receiving the plurality of speech recognition results from the speech recognition module comprises receiving the plurality of speech recognition results from the speech recognition module located on the server.

(121) Example 56 includes the subject matter of any of Examples 54 and 55, and wherein distorting the audio input comprises removing an internal segment of the audio input.

(122) Example 57 includes the subject matter of any of Examples 54-56, and wherein removing the internal segment of the audio input comprises removing a segment of the audio input having an amplitude with a predefined relationship to an amplitude threshold.

(123) Example 58 includes the subject matter of any of Examples 54-57, and wherein removing the internal segment comprises removing the segment having an amplitude below the amplitude threshold.

(124) Example 59 includes the subject matter of any of Examples 54-58, and wherein distorting the audio input comprises expanding a length of a segment of the audio input having an amplitude with a predefined relationship to an amplitude threshold.

(125) Example 60 includes the subject matter of any of Examples 54-59, and wherein expanding the length of the segment comprises expanding a length of the segment having amplitude below the amplitude threshold.

(126) Example 61 includes the subject matter of any of Examples 54-60, and wherein distorting the audio input comprises inserting a pause at a phonetic split point of the audio input identified by performing speech recognition on the audio input.

(127) Example 62 includes the subject matter of any of Examples 54-61, and wherein distorting the audio input comprises performing at least one of: (i) speeding up the audio input, (ii) slowing down

the audio input, (iii) adjusting a pitch of the audio input, or (iv) introducing noise to the audio input.

(128) Example 63 includes the subject matter of any of Examples 54-62, and wherein performing speech recognition on the audio input and the distorted audio variations to produce the plurality of speech recognition results comprises performing speech recognition on the audio input and the distorted audio variations to produce at least one hundred speech recognition results.

(129) Example 64 includes a method for speech recognition on a speech recognition server, the method comprising: receiving, on the speech recognition server, audio input captured by a computing device; distorting, on the speech recognition server, the audio input to produce a plurality of distorted audio variations; performing, on the speech recognition server, speech recognition on the audio input and each of the distorted audio variations to produce a plurality of speech recognition results; and transmitting the plurality of speech recognition results to the computing device.

(130) Example 65 includes the subject matter of Example 64, and wherein distorting the audio input comprises removing an internal segment of the audio input.

(131) Example 66 includes the subject matter of any of Examples 64 and 65, and wherein removing the internal segment of the audio input comprises removing a segment of the audio input having an amplitude with a predefined relationship to an amplitude threshold.

(132) Example 67 includes the subject matter of any of Examples 64-66, and wherein removing the internal segment comprises removing the segment having an amplitude below the amplitude threshold.

(133) Example 68 includes the subject matter of any of Examples 64-67, and wherein distorting the audio input comprises expanding a length of a segment of the audio input having an amplitude with a predefined relationship to an amplitude threshold.

(134) Example 69 includes the subject matter of any of Examples 64-68, and wherein expanding the length of the segment comprises expanding a length of the segment having amplitude below the amplitude threshold.

(135) Example 70 includes the subject matter of any of Examples 64-69, and wherein distorting the audio input comprises inserting a pause at a phonetic split point of the audio input identified by performing speech recognition on the audio input.

(136) Example 71 includes the subject matter of any of Examples 64-70, and wherein distorting the audio input comprises performing at least one of: (i) speeding up the audio input, (ii) slowing down the audio input, (iii) adjusting a pitch of the audio input, or (iv) introducing noise to the audio input.

(137) Example 72 includes the subject matter of any of Examples 64-71, and wherein performing speech recognition on the audio input and the distorted audio variations to produce the plurality of speech recognition results comprises performing speech recognition on the audio input and the distorted audio variations to produce at least one hundred speech recognition results.

(138) Example 73 includes a method for interaction with a virtual personal assistant on a computing device, the method comprising: displaying an avatar of the virtual personal assistant in a disengaged state on a display of the computing device; determining, on the computing device, whether a user of the computing device has requested activation of the virtual personal assistant; displaying, on the computing device, the avatar in a ready state in response to determining that the user has requested activation of the virtual personal assistant; determining, on the computing device, an engagement level of the user while the avatar is in the ready state; displaying, on the computing device, the avatar in an engaged state in response to the user having an engagement level greater than a threshold level; and displaying, on the computing device, the avatar in the disengaged state in response to the user having an engagement level less than the threshold level.

(139) Example 74 includes the subject matter of Example 73, and further comprising monitoring the engagement level of the user while the avatar is in the engaged state.

(140) Example 75 includes the subject matter of any of Examples 73 and 74, and wherein determining whether the user has requested activation of the virtual personal assistant comprises: receiving eye tracking data from an eye tracking sensor of the computing device; and determining whether the user has focused on the avatar based on the eye tracking data.

(141) Example 76 includes the subject matter of any of Examples 73-75, and wherein determining whether the user has requested activation of the virtual personal assistant comprises: receiving audio input from an audio sensor of the computing device; and detecting a code word uttered by the user based on the audio input.

(142) Example 77 includes the subject matter of any of Examples 73-76, and wherein determining the engagement level of the user comprises: receiving eye tracking data from an eye tracking sensor of the computing device; determining, based on the eye tracking data, whether the user has visually focused on the avatar for a period of time greater than a threshold time; determining that the user has an engagement level greater than the threshold level in response to determining that the user has visually focused on the avatar for a period of time greater than the threshold time; and determining that the user has an engagement level less than the threshold level in response to determining that the user has not visually focused on the avatar for a period of time greater than the threshold time.

(143) Example 78 includes the subject matter of any of Examples 73-77, and wherein determining whether the user has visually focused on the avatar for a period of time greater than the threshold time comprises ignoring glances away from the avatar for a second period of time less than a second threshold time.

(144) Example 79 includes the subject matter of any of Examples 73-78, and wherein determining the engagement level of the user comprises receiving audio input from an audio sensor of the computing device.

(145) Example 80 includes the subject matter of any of Examples 73-79, and wherein: displaying the avatar in the ready state comprises displaying an anthropomorphic representation of eye contact of the avatar with the user; and displaying the avatar in the engaged state comprises displaying an anthropomorphic representation of eye contact with the user.

(146) Example 81 includes the subject matter of any of Examples 73-80, and wherein displaying the anthropomorphic representation of eye contact comprises: receiving at least one of eye tracking data or head position data from an eye tracking sensor of the computing device; and displaying an anthropomorphic representation of eyes of the avatar following the user.

(147) Example 82 includes the subject matter of any of Examples 73-81, and wherein displaying the avatar of the virtual personal assistant in the disengaged state comprises: displaying the avatar as semitransparent, allowing a user interface element of the computing device to shine through the avatar; displaying the avatar in the ready state comprises decreasing a transparency of the avatar; and displaying the avatar in the engaged state comprises one of: decreasing the transparency of the avatar or eliminating the transparency of the avatar.

(148) Example 83 includes the subject matter of any of Examples 73-82, and wherein: displaying the avatar of the virtual personal assistant in the disengaged state comprises displaying the avatar at a position on the display away from an active application of the computing device; displaying the avatar in the ready state comprises moving the avatar on the display to a position closer to an active application of the computing device; and displaying the avatar in the engaged state comprises moving the avatar on the display to a position over an active application of the computing device without preventing input from the user to the active application.

(149) Example 84 includes the subject matter of any of Examples 73-83, and wherein: displaying the avatar in the ready state comprises moving the avatar on the display to a position closer to a position where the user is focused, based on eye tracking data received from an eye tracking sensor of the computing device; and displaying the avatar in the engaged state comprises moving the avatar on the display to a position closer to a position where the user is focused, based on eye

tracking data received from an eye tracking sensor of the computing device.

(150) Example 85 includes the subject matter of any of Examples 73-84, and wherein: displaying the avatar of the virtual personal assistant in the disengaged state comprises displaying the avatar as invisible; and displaying the avatar in the ready state comprises displaying the avatar as visible.

(151) Example 86 includes the subject matter of any of Examples 73-85, and further comprising: performing speech recognition while the avatar is in the engaged state; determining whether the user has issued a command based on the speech recognition; and displaying the avatar in a working state in response to determining the user has issued the command.

(152) Example 87 includes the subject matter of any of Examples 73-86, and wherein displaying the avatar in the working state comprises at least one of: (i) increasing a transparency of the avatar or (ii) moving the avatar on the display to a position away from an active application of the computing device.

(153) Example 88 includes a computing device comprising: a processor; and a memory having stored therein a plurality of instructions that when executed by the processor cause the computing device to perform the method of any of Examples 43-87.

(154) Example 89 includes one or more machine readable storage media comprising a plurality of instructions stored thereon that in response to being executed result in a computing device performing the method of any of Examples 43-87.

(155) Example 90 includes a computing device comprising means for performing the method of any of Examples 43-87.

Claims

1. An electronic user device comprising: a screen; machine-readable instructions; and at least one processor circuit to be programmed by the machine-readable instructions to at least: cause the screen to present an avatar in a first position, the avatar having eyes, a gaze of the eyes of the avatar in a first direction relative to a user; detect activation of an application installed on the electronic user device; and cause the screen to present the avatar in a second position relative to an image associated with the application, the gaze of the eyes of the avatar in a second direction relative to the user, the second direction different than the first direction.
2. The electronic user device of claim 1, wherein a face of the avatar includes human facial features.
3. The electronic user device of claim 1, wherein one or more of the at least one processor circuit is to cause the screen to present the avatar as facing toward the user in the first position.
4. The electronic user device of claim 1, wherein the gaze of the eyes of the avatar in the second direction is directed away from the user.
5. The electronic user device of claim 1, wherein one or more of the at least one processor circuit is to cause the screen to present the avatar spaced apart from the image.
6. The electronic user device of claim 1, wherein one or more of the at least one processor circuit is to cause the screen to present content with the avatar when the avatar is in the first position.
7. The electronic user device of claim 6, wherein one or more of the at least one processor circuit is to cause the screen to present the avatar in closer proximity to the content when the avatar is in the first position than to the image when the avatar is in the second position.
8. The electronic user device of claim 1, further including a microphone.
9. At least one non-transitory machine-readable storage medium comprising machine-readable instructions to cause at least one processor circuit of an electronic user device to at least: cause a screen of the electronic user device to present an avatar, the avatar in a first position, the avatar having eyes, a gaze of the eyes of the avatar in a first direction relative to a user; detect a user input associated with an application installed on the electronic user device; and responsive to the user input, cause the screen to present the avatar in a second position relative to content associated with

- the application, the gaze of the eyes of the avatar in a second direction relative to the user, the second direction different than the first direction.
10. The at least one non-transitory machine-readable storage medium of claim 9, wherein the avatar has a face including human features.
 11. The at least one non-transitory machine-readable storage medium of claim 9, wherein the machine-readable instructions are to cause one or more of the at least one processor circuit to cause the screen to present other content with the avatar when the avatar is in the first position.
 12. The at least one non-transitory machine-readable storage medium of claim 11, wherein a face of the avatar is opaque.
 13. The at least one non-transitory machine-readable storage medium of claim 9, wherein the gaze of the eyes of the avatar in the second direction is directed away from the user.
 14. The at least one non-transitory machine-readable storage medium of claim 9, wherein one or more of the at least one processor circuit is to cause the screen to present the avatar as facing out of the screen in the first position.
 15. An electronic user device comprising: a screen; machine-readable instructions; at least one processor circuit to be programmed by the machine-readable instructions to: cause the screen to present an avatar in a first position relative to content associated with a first application installed on the electronic user device; detect a user input associated with a second application installed on the electronic user device; and responsive to the user input, cause the screen to present the avatar in a second position relative to content associated with the second application, the first position of the avatar in closer proximity to the content associated with the first application than the proximity of the avatar to the content associated with the second application in the second position.
 16. The electronic user device of claim 15, wherein the avatar has eyes and one or more of the at least one processor circuit is to: cause the screen to present a gaze of the eyes of the avatar in a first direction relative to a user when the avatar is in the first position; and cause the screen to present the gaze of the eyes of the avatar in a second direction relative to the user when the avatar is in the second position, the second direction different than the first position.
 17. The electronic user device of claim 16, wherein one or more of the at least one processor circuit is to cause the screen to present the gaze of the eyes of the avatar toward the user in the first direction.
 18. The electronic user device of claim 15, further including a sensor to track an eye of a user.
 19. The electronic user device of claim 15, wherein the avatar has a human face appearance.
 20. The electronic user device of claim 19, wherein the face of the avatar is opaque.
-