

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12399600
Kind Code	B1
Date of Patent	August 26, 2025
Inventor(s)	Shin; Dongeek et al.

Contact space determination for user interface control

Abstract

Described techniques enable a virtual touchpad for a wearable device by receiving, in response to a touch event, a sensor signal from a motion sensor coupled to the wearable device, the wearable device providing a user interface with a rendering space that is defined relative to the wearable device. The touch event may be associated with a contact space of a contact map defined with respect to the wearable device, based on the sensor signal, and a user interface event of the user interface that corresponds to the contact space may be rendered within the rendering space.

Inventors: Shin; Dongeek (San Jose, CA), Kulkarni; Anuva (Santa Clara, CA), Xu; Tianyu (Mountain View, CA), Mu; Lidan (Mountain View, CA)

Applicant: Google LLC (Mountain View, CA)

Family ID: 1000007709242

Assignee: Google LLC (Mountain View, CA)

Appl. No.: 18/586122

Filed: February 23, 2024

Publication Classification

Int. Cl.: G06F3/04815 (20220101); G06F3/01 (20060101); G06F3/0346 (20130101); G06F3/04845 (20220101); G06F3/0487 (20130101)

U.S. Cl.:

CPC G06F3/04815 (20130101); G06F3/011 (20130101); G06F3/017 (20130101); G06F3/0346 (20130101); G06F3/04845 (20130101); G06F3/0487 (20130101);

Field of Classification Search

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
10181219	12/2018	Murphy	N/A	A63F 13/26
11416075	12/2021	Gong et al.	N/A	N/A
12086919	12/2023	Zhang	N/A	G06F 3/013
2016/0267712	12/2015	Nartker	N/A	G06F 3/0325
2018/0005439	12/2017	Evans	N/A	G06T 19/006
2018/0322706	12/2017	Drouin	N/A	G06F 3/011
2020/0104039	12/2019	Robertson et al.	N/A	N/A
2020/0278763	12/2019	Munakata	N/A	G06F 3/0414
2021/0043011	12/2020	Gates	N/A	G02B 27/017
2021/0056764	12/2020	Lacey et al.	N/A	N/A
2022/0067378	12/2021	Rabinovich	N/A	G06F 3/011
2022/0308655	12/2021	Zhu et al.	N/A	N/A
2023/0117197	12/2022	Stolzenberg	N/A	N/A

OTHER PUBLICATIONS

International Search Report and Written Opinion for PCT Application No. PCT/US2025/016441, mailed on Jun. 3, 2025, 12 pages. cited by applicant

Primary Examiner: Lin; Chun-Nan

Attorney, Agent or Firm: Brake Hughes Bellermann LLP

Background/Summary

BACKGROUND

(1) Extended Reality (XR) devices typically provide a user interface (UI). Various types of UI control exist that are compatible with UIs of XR devices. For example, UI control may be implemented using physical buttons, image-based gesture recognition, or through the use of separate devices (e.g., handheld controllers).

SUMMARY

(2) As described herein, a computing device may be configured to use touch events at one or more surfaces of the computing device to implement user interface control. Specifically, as described herein, a computing device such as an XR device may be configured to use touch events at one or more surfaces of the XR device to implement user interface control. The user interface control may include any manipulation of an XR user interface element or aspect, including spatial rendering of desired effects. For example, a location or position on an XR device at which a touch event occurs may be related to a rendering event at a corresponding rendering location within the user interface of the XR device. In other examples, a contact space that occurs in conjunction with a touch event may be captured, where a contact space refers to, e.g., a size, area, shape, or location of a surface area of an XR device surface at which contact of a touch event occurs, including movement of the touch event across the surface area of the XR device. A rendering event may be initiated in

response to characteristics (e.g., a location, shape, size, area, or movement) of the contact space. Thus, XR devices may be provided with a highly customizable set of user interface controls, which may be convenient to use. Moreover, such user interface controls may be provided without consuming dedicated space on the surface of the XR device (such as might be required for installing buttons or other control devices). Although at least some of the description is related to an XR device, the concepts described herein can be applied to any type of computing device such as a mobile phone, laptop computing device, smart device, and/or so forth.

(3) In a general aspect, a method includes receiving, in response to a touch event, a sensor signal from a motion sensor mounted on a device, the device providing a user interface with a rendering space that is defined relative to the device. The method includes relating the touch event to a contact space defined with respect to the device, based on the sensor signal, and rendering a user interface event of the user interface within the rendering space that corresponds to the contact space.

(4) In another general aspect, a computer program product is tangibly embodied on a non-transitory computer-readable storage medium and comprises instructions. When executed by at least one computing device (e.g., by at least one processor of the at least one computing device), the instructions are configured to cause the at least one computing device to receive, in response to a touch event, a sensor signal from a motion sensor mounted on a device, the device providing a user interface with a rendering space that is defined relative to the device. When executed by at least one computing device (e.g., by at least one processor of the at least one computing device), the instructions are configured to cause the at least one computing device to relate the touch event to a contact space defined with respect to the device, based on the sensor signal. When executed by at least one computing device (e.g., by at least one processor of the at least one computing device), the instructions are configured to cause the at least one computing device to render a user interface event of the user interface within the rendering space that corresponds to the contact space.

(5) In another general aspect, a wearable device includes at least one frame for positioning the wearable device on a body of a user, at least one display, at least one processor, and at least one memory storing instructions. When executed, the instructions cause the at least one processor to receive, in response to a touch event, a sensor signal from a motion sensor mounted on a device, the device providing a user interface with a rendering space that is defined relative to the device. When executed, the instructions cause the at least one processor to relate the touch event to a contact space defined with respect to the device, based on the sensor signal. When executed, the instructions cause the at least one processor to render a user interface event of the user interface within the rendering space that corresponds to the contact space.

(6) The details of one or more implementations are set forth in the accompanying drawings and the description below. Other features will be apparent from the description and drawings, and from the claims.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1A illustrates an example implementation of a system for contact space determination for user interface (UI) control.

(2) FIG. 1B illustrates a block diagram for implementing the system for contact space determination for user interface control of FIG. 1A.

(3) FIG. 1C illustrates additional use case scenarios for implementing the system of FIG. 1B.

(4) FIG. 2 is a flowchart illustrating example operations of the system of FIG. 1B.

(5) FIG. 3 is a block diagram of an alternate example of the system of FIG. 1B.

(6) FIG. 4 is a block diagram of an alternate example of the system of FIG. 1B.

- (7) FIG. 5 illustrates an example use case scenario for the systems of FIGS. 1B and 4.
- (8) FIG. 6 illustrates a second example use case scenario for the systems of FIGS. 1B and 4.
- (9) FIG. 7 is a third person view of a user in an ambient computing environment.
- (10) FIGS. 8A and 8B illustrate front and rear views of an example implementation of a pair of smartglasses.

DETAILED DESCRIPTION

(11) Described systems and techniques enable use of touch events occurring at a computing device for user interface (UI) control of a user interface of the computing device. For example, the systems and techniques described herein enable use of touch events occurring at a computing device such as an XR device for user interface control of a user interface of the XR device. For example, such touch events may include any contacts between the XR device, e.g., a surface of the XR device, and any other object(s), including, e.g., a finger or a stylus. Touch events may include, e.g., tapping, pressing, holding, swiping, pinching, zooming, rotating, or any other type of touch or contact.

(12) Such touch events may be used to control rendering events or other user interface events occurring within a rendering space of the user interface of the XR device. For example, one or more motion sensors of the XR device may be used to define, characterize, or otherwise determine a contact space that refers to a portion of the XR device at which contact occurs, as in the case of a touch event.

(13) The contact space may be characterized with respect to a size, shape, area, location, or movement of a contact or other touch event at or across the XR device surface. For example, the contact space may be determined at a certain position(s) and with a certain size/shape within a contact map. The contact map may provide a bounded area of the XR device surface in which touch events and corresponding contact spaces may be detected.

(14) The contact space may be used to control the user interface of the XR device, including, e.g., initiating a rendering event at a location within the rendering space that corresponds to the location of the contact space within the contact map. In the present description, as described in more detail, below, a rendering space refers to any virtual or physical space around an XR device in which user interface elements or aspects may be rendered. For example, a rendering space may include a virtual reality room or other setting rendered by the XR device. In an augmented reality context, the rendering space may refer to a physical area surrounding the XR device in which virtual objects or other interface elements may be rendered. Thus, a rendering event includes any user interface event that is caused to occur by the XR device within the rendering space of the user interface of the XR device. Although at least some of the description is related to an XR device, the concepts described herein can be applied to any type of computing device such as a mobile phone, laptop computing device, smart device, and/or so forth.

(15) For example, two fingers touching the contact map may occupy more contact space than one finger touching the contact map. In either case, the contact space may be mapped to the rendering space, so that a subsequent rendering event or other user interface event may occur in a defined manner within the rendering space. For example, when creating a punch-through effect for a virtual reality (VR) context, a two-finger contact space may result in a larger punch-through area than a one-finger contact space.

(16) At least one technical problem solved by the described techniques includes providing user interface control in XR environments. At least one technical problem solved by the described techniques includes providing user interface control for XR devices that have a small form factor and correspondingly small device area to use for providing user interface control. At least one technical problem solved by the described techniques includes providing a large number of configurable user interface control options that can be related to a user interface of XR devices.

(17) At least one solution to the above and other technical problems includes defining a contact map with respect to a device area, e.g., device surface area, of an XR device. Then, by detecting a

contact space associated with a touch event that occurs within the contact map, the contact space may be related to a rendering space of a user interface of the XR device. Since virtually any contact space may be detected at any defined location of the device surface area/contact map, corresponding user interface control options (e.g., a large number of corresponding user interface control options) may be made available.

(18) In existing XR devices, as referenced above, user interface control may be achieved using a variety of known techniques. For example, buttons (or other physical input devices) may be installed on a head mounted device (HMD), such as glasses or goggles, and user interface events may be initiated in response to operation of the physical buttons. However, such physical input devices require hardwiring and occupy space within a surface area of an XR device, thereby potentially increasing (or preventing a reduction of) a form factor of the XR device.

(19) In other examples, separate devices may be in communications with an XR device, such as a HMD, and may be used for user interface control of the XR device. For example, a smartwatch, smartphone, smartring, or handheld controller may be used to control an HMD. However, some HMDs do not provide or interact with such separate devices in a desired manner, and not all users may have access to such separate devices at a time of use of an HMD. Moreover, it may be difficult to configure desired types of interoperability, with desired levels of accuracy, between the HMD and the various control devices.

(20) In still other examples, a sensor of the XR device may be used to capture sensor signals that reflect interactions of a user with the XR device. For example, a camera of an HMD may be used to provide eye gaze tracking, or a motion sensor may be used to implement head gaze tracking.

(21) In other examples, a camera may be used to capture an image of a hand gesture of a user, or a motion sensor may be used to capture a touch event of a user with the XR device. Then, a machine learning (ML) model may be used to classify the sensor signals as corresponding to one or more defined gestures, which may then be related to implementing corresponding user interface events. For example, a convolutional neural network (CNN) may be trained to recognize a defined set of hand gestures (e.g., pinching, waving, or pointing), each of which may then be related to corresponding user interface events (e.g., zooming, scrolling, or selecting a user interface element).

(22) In related examples, a touch event (e.g., tapping) may be recognized by using a motion sensor mounted on an HMD. For example, a linear regression model may be trained to relate one or more touch events to one or more classifications of gestures for user interface control. Similar to image-based gesture control, such techniques are generally limited to a small set of available, recognizable gestures. Moreover, such techniques may be limited to XR device surfaces that are proximate to the motion sensor being used for signal detection. Still further, such techniques, and various other ones of the user interface control techniques referenced above, may be difficult to calibrate across multiple XR devices or types of XR devices, so as to be functional in a desired manner for a large number and variety of users.

(23) In example techniques described herein, a contact map may be defined with respect to any desired surface of an XR device. For example, a contact map may be defined with respect to an arm of a pair of smartglasses, or a face of a pair of XR goggles. Within the contact map, any device surface portion that is touched may be associated with a contact space. Characteristics of the contact space may then be used to implement user interface control for the XR device being used.

(24) Using such techniques, a location of a contact map and/or contact space may be mapped to a rendering space of the XR user interface. That is, extended Reality (XR) devices are capable of spatial rendering, in which a user interface (UI) is defined relative to a corresponding XR device, and user interface events are rendered within a rendering space that is defined relative to a physical position of the corresponding XR device. For example, a user of XR goggles may experience a virtual object rendered at a certain distance from the XR goggles, or at a certain height or other position defined relative to the XR goggles.

(25) Therefore, for example, a contact space on one side of the XR device may correspond to a

rendering or other user interface control event on the same side of the rendering space of the XR device. For example, touching the right side of the XR device may result in a control event at a right side of the rendering space. In other examples, as referenced above, a size of the contact space may dictate a size of a corresponding rendering space or rendering event.

(26) Detected contact spaces need not be limited to a pre-defined set of detectable gestures for user interface control. Rather, virtually any desired gesture may be defined and detected with respect to a corresponding contact space. Moreover, described techniques may be easily calibrated across many types of computing devices such as XR devices.

(27) FIG. 1A illustrates an example implementation of a system for contact space determination for user interface control. In the example of FIG. 1, a wearable device **102** is illustrated as being worn by a user **104**. For example, the wearable device **102** may represent a head-mounted device (HMD), such as smartglasses or smart goggles. Although this example is directed to a HMD device, the concepts described herein can be applied to any type of computing device such as a mobile phone, laptop computing device, smart device, and/or so forth.

(28) FIG. 1A further illustrates a device surface **102a** of the HMD **102** and a user interface **102b** provided by the HMD **102**. For example, the device surface **102a** may include any one or more surface areas of the HMD **102** that may be touched or otherwise contacted by the user **104**. For example, when the HMD **102** represents a pair of smart goggles, the device surface **102a** may represent a front surface of the smart goggles (e.g., viewable from a point of view of facing the user **104** while the user **104** is wearing the smart goggles).

(29) More generally, the device surface **102a** may represent a top, side, or bottom surface of smart goggles. When the HMD **102** represents smartglasses, as in FIGS. 8A and 8B, the device surface **102a** may represent any surface on the arms or frame of the smart glasses. Although shown in the singular in the simplified example of FIG. 1A, it will be appreciated that two or more device surfaces may be defined for purposes of providing user interface control as described herein.

(30) The user interface **102b** is illustrated from a point of view of the user **104** when wearing the HMD **102**. In other words, the user interface **102b** is seen by the user **104** when wearing the HMD **102**. The user interface **102b** may include, for example, a Virtual Reality (VR) interface in which the user **104** is immersed in a 360° field of view provided by an associated application. In other examples, the user interface **102b** may include an Augmented Reality (AR) interface in which the user **104** has visibility with respect to a surrounding physical environment, and various virtual user interface elements or aspects are superimposed on, or otherwise added to or shown in the context of, the physical environment. Any other suitable type of user interface that may be provided by the HMD **102** may be used, as well.

(31) The HMD **102** may include one or more motion sensors, represented in FIG. 1A in block diagram form by a motion sensor **106** and a motion sensor **108**. For example, the motion sensors **106**, **108** may represent one or more of an inertial measurement unit (IMU), an accelerometer, a gyroscope, a magnetometer, or any other known or future device for detecting a motion of the HMD **102**.

(32) As described in detail, below, sensor signals from one or more of the motion sensors **106**, **108** may be processed by the HMD **102** (and/or connected devices) to detect and process user interactions with a contact map **110** defined with respect to the device surface **102a**. Although illustrated in the singular in the simplified example of FIG. 1A, the contact map **110** may represent one or more defined areas of the device surface **102a**.

(33) The contact map **110** may be used to control operations of the user interface **102b**. For example, the user interface **102b** is illustrated as including a rendering space **112**, which may include a spatially defined region or area of the user interface **102b** and/or of a surrounding physical environment of the user **104**. For example, in a VR context, the rendering space **112** may include an immersive, rendered 360° VR environment. In an AR context, the rendering space **112** may include a virtual spatial region overlaid on a surrounding physical environment of the user

104. In this way, for example, virtual objects or other virtual user interface elements (e.g., user interface control elements) may be placed within the rendering space **112**.

(34) The rendering space **112** may thus represent a coordinate space of any physical or virtual location, or combination thereof, that may be defined with respect to the HMD **102**. For example, the rendering space may represent coordinates within a physical space such as a room, an arena, or an outdoor space, or within a virtual space, which could include virtually any virtual location capable of being rendered by the HMD **102**. The rendering space **112** may also represent a common coordinate system that is applied across both a physical and virtual environment.

(35) Within the defined area of the contact map **110**, a touch event **114** may be associated with a contact space **116** for control of operations of the user interface **102b**, including control of events within the rendering space **112**. That is, using techniques described below, sensor signals from the motion sensors **106**, **108** that occur in response to the touch event **114** may be processed to determine the contact space **116** within the contact map **110**.

(36) For example, the user **104** may touch the contact map **110** with one or more fingers, or with a stylus, or using any other suitable object to physically interact with or touch the contact map **110**. In many of the following examples, for the sake of simplicity and consistency of explanation, the touch event **114** (and other touch events) are generally referenced with respect to a finger(s) of the user **104** touching the contact map **110**.

(37) The contact space **116** thus refers to an included area or subset of space of the contact map **110**, which corresponds in size, extent, position, and direction to the touch event **114**. For example, a touch event that consumes more area of the contact map **110** will result in a larger contact space.

(38) For example, in FIG. 1A, a touch event **118** is illustrated as being associated with a larger contact space **120** than the contact space **116** associated with the touch event **114**. For example, the touch event **118** may relate to a greater extent or portion of a finger of the user **104** being positioned on the contact map **110** than the touch event **114**. For example, the touch event **114** may correspond to a fingertip of the user **104** being in contact with the contact map **110**, while the touch event **118** may correspond to a larger portion of the finger of the user **104** being in contact with the contact map **110**.

(39) Further, and as also illustrated, a touch event, such as the touch event **118**, need not be static, and may reflect or include a dynamic movement of a finger of the user **104** within, and with respect to, the contact map **110**. Additionally, the touch event **114** and the touch event **118** may occur separately or together, and various other types of touch events may be detected, as well, some of which are described below by way of further example.

(40) In response to the contact space **116** being detected, a corresponding user interface event **116a** may be implemented within a context of the user interface **102b**, e.g., within the rendering space **112**. Similarly, in response to the contact space **120** being detected, a corresponding user interface event **120a** may be implemented.

(41) The user interface events **116a**, **120a** may represent virtually any user interface event that is associated with operation or use of the user interface **102b**. Such user interface events may include, by way of non-limiting example, user interface object or element selections, or any user interface action associated with operations of an application that provides the user interface **102b** (e.g., a browser, a game, an image/video display, or any other type of application).

(42) The user interface events **116a**, **120a** may also represent any event(s) that affects operations of the user interface **102b** that may be common to the HMD **102** across multiple applications or uses of the HMD **102**. For example, as described in more detail with respect to FIGS. 5 and 6, the user interface event **116a** may represent a punch-through event or punch-through window, in which a window or area is provided within an immersive VR environment to enable the user **104** to easily view a surrounding physical environment.

(43) Advantageously in FIG. 1A, a surface area of the contact map **110** may be mapped in virtually any desired relationship to the user interface **102b**, e.g., to the rendering space **112**. Thus, for

example, a position of any touch event and associated contact space within the contact map **110** may be spatially mapped to the user interface **102b**, e.g., to the rendering space **112**.

(44) For example, as mentioned above, the device surface **102a** may represent a forward-facing surface of a pair of smart goggles as the HMD **102**. In such examples, the touch event **114** may represent a finger of the user **104** touching the contact map **110** in a direction to the right side of the user **104**, while the touch event **118** may represent a finger of the user **104** touching the contact map **110** in a direction to the left side of the user **104** (that is, reversed from the illustrated view of FIG. 1A, in which the touch event **114** is illustrated to the left of the touch event **118**).

(45) Then, the contact space **116** may be mapped to the rendering space **112** with the user interface event **116a** being implemented toward the right side of the user **104** as the user **104** looks ahead, and, similarly, with the user interface event **120a** being implemented toward the left side of the user **104**. For example, if the user interface events **116a**, **120a** include punch-through events, then the user **104** may easily obtain a desired punch-through event (e.g., may obtain visibility in a desired direction(s)), simply by touching the contact map **110** at corresponding locations.

(46) More generally, virtually any desired mapping may be implemented between the contact map **110** and the user interface **102b**, e.g., within the rendering space **112**. Such mappings may be implemented regardless of where on the device surface **102a** the contact map **110** is positioned.

(47) Moreover, such mappings may be implemented in 1, 2, or 3 dimensions within the rendering space **112**. For example, the touch event **114** and contact space **116** may be related to a defined point within the rendering space **112**, or movement of the contact space **120** across a surface of the contact map **110** may be related to x, y coordinates within the rendering space **112**.

(48) Further, the contact map **110** may be positioned on a portion of the device surface **102a** that is positioned in a direction that is perpendicular to an x,y plane of the rendering space **112**. For example, the HMD **102** may include smartglasses, and the device surface **102a** may represent an arm of the smartglasses supported on an ear of the user **104** and perpendicular to lenses of the smartglasses.

(49) Then, the touch event **118** may represent a motion occurring along a length of the arm of the smartglasses, in a z or depth direction within the rendering space **112**. For example, the rendering space **112** may include the type of immersive 3D VR environment referenced above, or the type of AR environment referenced above that is rendered within a 3D environment of the user **104**. For example, a 3D VR environment may include virtual objects positioned at various depths, while the AR environment may include virtual objects placed relative to real-world objects in a surrounding environment of the user **104**.

(50) Therefore, using described techniques, the user **104** may easily select or utilize any such virtual objects, or reference any real-world objects, at any desired depth. Moreover, multiple implementations of the contact map **110** may be included on the HMD, to provide full 3D control within the rendering space **112**.

(51) For example, a first implementation of the contact map **110** may be placed on a front or left surface of the HMD **102**, while a second implementation of the contact map **110** may be placed on a right surface of the HMD **102**. The first implementation of the contact map **110** may then be used to control left-right movements, e.g., of a cursor in the rendering space **112**, while the second implementation of the contact map **110** may be used to control forward-back movement, e.g., of the cursor. Such embodiments therefore enable easy and intuitive user interface control, even in 3D immersive environments.

(52) Various other types of user interface control may be provided, as well. For example, the HMD **102** may include various types of hardware sensors or other devices, including output devices. For example, the HMD **102** may include a camera(s), microphone(s), speaker(s), and various other hardware elements. Therefore, the contact map **110** may enable control of corresponding user interfaces with such devices, in order to control their operation(s). Using techniques described herein, such control may include directional aspects.

(53) For example, multiple cameras, speakers, or microphones may be positioned around a frame of the HMD **102**, each being, e.g., proximate to a corresponding one of the motion sensors **106**, **108**. Then, for example, the user **104** may preferentially select and operate a desired one of such devices by touching a corresponding portion of the contact map **110**. For example, the user **104** may control operations of a microphone or speaker next to the motion sensor **108** by virtue of executing the touch event **118** to create the contact space **120** in a vicinity of the motion sensor **108**.

(54) More generally, as described in detail with respect to FIGS. **4-6**, placement of multiple motion sensors **106**, **108** relative to the device surface **102a** results in a single touch event within the contact map **110** causing a distributed profile of sensor signals that facilitates determination of a corresponding contact space. For example, as shown with respect to FIGS. **5** and **6**, the touch event **118** will result in a first sensor signal distribution in which a sensor signal of the motion sensor **108** is greater than a sensor signal of the motion sensor **106**, while the touch event **114** will result in a second sensor signal distribution in which a sensor signal of the motion sensor **106** is greater than a sensor signal of the motion sensor **108**. Such distributions of sensor signals may thus be used to infer or otherwise determine a location, movement, position, extent, and nature of the contact spaces **116**, **120** within the contact map **110**, which, as already described, may be used to provide positional and directional user interface control with respect to the user interface **102b** and the rendering space **112**.

(55) Thus, the contact spaces **116**, **120** may provide information for a total surface area of the contact map that is consumed by the corresponding touch events **114**, **118**, and/or a location within the contact map **110** that may be related to a corresponding location within the rendering space **112**, or may otherwise be used to control one or more user interface event(s) **120a**, **116a**.

(56) Described techniques may thus provide user interface control over a large field of view (FOV) while providing a complex and highly customizable gesture set. Moreover, even with respect to a single gesture, described techniques may enable differences in corresponding user interface events that depend on a recognized location within the contact map **110**, and/or on a location of one or more of the contact map(s) **110** on the device surface **102a**.

(57) For example, as shown in the block diagram of FIG. **1B**, the HMD **102** may include a processor **122** (which may represent one or more processors), as well as a memory **124** (which may represent one or more memories (e.g., non-transitory computer readable storage media)). Although not shown separately in FIG. **1B**, the HMD **102** may also include a battery, which may be used to power operations of the processor **122**, the memory **124**, and various other resources of the HMD **102**. More detailed examples of the HMD **102** and various associated hardware/software resources are provided below, e.g., with respect to FIGS. **7**, **8A**, and **8B**.

(58) For purposes of the simplified example of FIG. **1B**, the HMD **102** should be further understood to include a depth map generator **126**, which may be configured to generate a depth map in conjunction with, or for use with, the rendering space **112**. The depth map generator **126** may utilize depth data from one or more passive or active depth sensor(s), not separately illustrated in FIG. **1B**, which may include, e.g., a time-of-flight (ToF) camera, LiDAR sensor, RGB camera, or GPS sensor. Using received, raw depth data, the depth map generator **126** may be configured to generate a depth map that captures information characterizing relative depths of detected or rendered objects with respect to a defined perspective or reference point. For example, such a depth map may be used to define or determine a depth of the user interface event **120a** or the user interface event **116a** within the rendering space **112**.

(59) The depth map generator **126** should be understood to represent and illustrate depth map software stored using the memory **124** and executed using the processor **122**, and configured to process depth-related data captured by one or more of the herein-referenced depth sensors, or other suitable depth sensor(s). The depth map generator **126** may thus be capable of determining a per-pixel depth of each pixel (and associated object or aspect) within the rendering space **112**. For example, such depth information may be captured and stored as a perspective image containing a

depth value instead of a color value in each pixel. A depth map may be generated and stored, e.g., as a 2D array, or as a depth mesh (e.g., a real-time triangulated mesh).

(60) A user interface controller **128** may be configured to provide and control the user interface **102b**, including the rendering space **112**. As described in more detail, below, with respect to the examples of FIGS. **3** and **4**, the user interface controller **128** may receive one or more sensor signals from the motion sensors **106**, **108**, in conjunction with a depth map from the depth map generator **126** and any other information used to provide the features and functions referenced above with respect to the contact map **110** in the context of FIG. **1A**.

(61) For example, the user interface controller may include a sensor signal preprocessor **130**. The sensor signal preprocessor **130** may be configured to buffer, synchronize, filter, debounce, or other preprocess sensor signals from the motion sensors **106**, **108**. For example, as the motion sensors **106**, **108** may be positioned at different positions in or on the HMD **102**, generated sensor signals may reach the user interface controller **128** at offset times, so that the sensor signal preprocessor **130** may be configured to buffer and synchronize the received signals. Debouncing may also be performed to avoid unwanted repetitions of received control signals.

(62) Further, the sensor signal preprocessor **130** may be configured to reject or filter unwanted or unrelated portions of received sensor signals. For example, the user **104** may simultaneously be exhibiting various head movements or other body movements not related to desired operations of the user interface **102b**, and such movements may be filtered by the sensor signal preprocessor **130**.

(63) In addition, the user interface controller **128** may represent or include software intended to be used with many different HMDs, or types of HMDs. Therefore, the sensor signal preprocessor **130** may be configured to provide calibration across all such HMDs. For example, different HMDs may have different types or numbers of motion sensors or other types of sensors, and/or may use different reference coordinate systems to define the rendering space **112** (e.g., may have different installed orientations). Consequently, as may be appreciated from the above description, the term calibration as used herein should be understood to refer to any processing actions taken to make a user experience and use of the user interface controller **128** consistent across multiple users and/or multiple HMDs.

(64) A gesture recognition model **132** may be configured to receive the preprocessed signals from the sensor signal preprocessor **130** and determine user actions (e.g., touch events **114**, **118**), contact spaces **116**, **118**, and any gestures and user interface control events (such as the user interface events **116a**, **120a**) associated therewith. For example, the gesture recognition model **132** may characterize above-referenced aspects of the contact spaces **116**, **120**, or other aspects, including, e.g., a total size, shape, movement, or duration thereof, including combinations of different, overlapping or successive touch events.

(65) It will be appreciated that the user interface controller **128** is highly configurable, and multiple types of known or future gesture recognition models may be used as the gesture recognition model **132** of FIG. **1B**. Known examples of gesture recognition models (e.g., CNN-based models), as referenced above, may be configured to classify a received sensor signal as corresponding to one of a pre-defined set of known gestures, which may then be used for certain types of user interface control. In contrast, in FIG. **1A** and FIG. **1B**, the gesture recognition model **132** may be configured to recognize the contact map **110** and the contact spaces **116**, **120** essentially as an image(s) constructed from the received and preprocessed sensor signals.

(66) Such images may be associated with classified or known gestures, but, in contrast with existing techniques, enable a wider range of such gestures, and combinations thereof, which are more easily implemented, customized, and configured than in existing techniques. Moreover, such recognized gestures may be implemented with a type and level of spatial awareness, e.g., within the rendering space **112**, not possible in existing techniques.

(67) For example, a mapping module **134** may be configured to perform image-to-image translation between outputs of the gesture recognition model **132** representing the contact map **110** and contact

space(s) **116**, **120** at a point(s) in time to the rendering space **112** of the user interface **102b**, as described above with respect to the user interface events **116a**, **120a**. Although illustrated separately in FIG. **1B** for purposes of explanation and illustration, it will be appreciated that the gesture recognition model **132** and the mapping module **134** may be implemented as a single neural network. For example, the mapping module **134** may be implemented as output layer(s) of the gesture recognition model **132** that replace or augment existing classification layers.

(68) Consequently, a rendering engine **136** may be configured to render the user interface **102b**, including the rendering space **112**. The rendering engine **136** may further render user interface events **116a**, **120a** at locations within the rendering space **112** that correspond to the contact spaces **116**, **120** within the contact map **110**. It will be appreciated that the rendering engine **136** may also implement virtually any features and functions associated with rendering the user interface **102b**, which may vary widely based on, e.g., a nature of an application or operating system of which the user interface controller **128** is a part, or operating in conjunction therewith. Such features/functions are not described herein, except as may be helpful in understanding operations of the user interface controller **128** in providing the contact map **110** and associated user interface control techniques.

(69) For example, by way of further illustration, FIG. **1C** illustrates additional use case scenarios enabled by described techniques. For example, in FIG. **1C**, a contact map **110c** is illustrated on a device surface **102c** of the HMD **102** that may be different than the device surface **102a** illustrated in FIG. **1A**. For example, the device surface **102c** may represent a side of the HMD **102**, or an arm of the HMD **102** when the HMD **102** is implemented using smartglasses.

(70) In FIG. **1C**, a touch event **138** may cause a contact space **140** to be detected, while a touch event **142** may cause a contact space **144** to be detected. The contact space **140** may be detected with respect to a user interface **140a** and the contact space **144** may be detected with respect to a user interface **144a**.

(71) For example, the rendering space **112** may be used to render two or more user interfaces in different portions or areas of the rendering space **112**. In FIG. **1C**, the user interface **140a** may represent an individual user interface with various components, or may represent an individual one of such components, such as a cursor or other control element. Similar comments apply to the user interface **144a**.

(72) In a first example use case illustrated with respect to the user interface **140a** and the user interface **144a**, the two contact spaces **140**, **144** may occur in a simultaneous or overlapping fashion. The user interface **140a** and the user interface **144a** may then also be presented in a corresponding simultaneous or overlapping fashion. For example, two cursors may be controlled at the same time.

(73) A second example use case is illustrated with respect to a user interface **144b**. In the example, the contact space **144** may correspond to the user interface **144b**, which may also represent an entire user interface or any portion thereof. As indicated by corresponding arrows in FIG. **1C**, a swiping touch gesture may be detected with respect to the touch event **142**, which may result in a corresponding movement of the user interface **144b** within the rendering space **112**, as illustrated.

(74) A third example use case is illustrated with respect to a user interface **144c**. In the example, the touch event **138** and the touch event **142** may occur at the same time and may be interpreted as corresponding a rendered size of the user interface **144c**. For example, a distance between the contact space **140** and the contact space **144** on the contact map **110c** may be interpreted as corresponding to a rendered size of the user interface **144c**.

(75) A fourth example use case may also be understood with respect to the user interface **144c**. In the example, touch-based shrinking or enlarging of the user interface **144c** may be implemented. For example, decreases in the distance between the contact space **140** and the contact space **144** may be interpreted as shrinking the user interface **144c**, while increases in the distance between the contact space **140** and the contact space **144** may be interpreted as enlarging the user interface

144c.

(76) A fifth example use case may also be understood with respect to the user interface **144c**. In the example, touch-based closing of the user interface **144c** may be provided using a variety of techniques. For example, a swiping gesture detected with respect to the contact space **144** may be interpreted as a command to close the user interface **144c**. In other examples, the contact space **144** may be used to control a cursor or other selection element that may be detected as selecting a closing action for the user interface **144c**.

(77) The various example use cases described herein should be understood to be non-limiting. For example, many other gestures and other uses of described contact spaces, and associated touch events, may be implemented for UI control. For example, touch events may be correlated with corresponding gestures and associated UI control based on an extent of pressure associated with a corresponding touch event, and/or on a duration of such a touch event.

(78) FIG. 2 is a flowchart illustrating example operations of the system of FIG. 1. In the example of FIG. 2, operations **202-206** are illustrated as separate, sequential operations. However, in various example implementations, the operations **202-206** may be implemented in a different order than illustrated, in an overlapping or parallel manner, and/or in a nested, iterative, looped, or branched fashion. Further, various operations or sub-operations may be included, omitted, or substituted.

(79) In FIG. 2, a sensor signal from a motion sensor coupled to a device may be received in response to a touch event, the device providing a user interface with a rendering space that is defined relative to the device (**202**). For example, sensor signal(s) from one or more of the motion sensors **106**, **108** may be received in response to one or both of the touch events **114**, **118**. In FIGS. **1A** and **1B**, description is provided with respect to the HMD **102** as an example device, and other example devices are provided below, with respect to FIG. 7. Any such device may provide the user interface **102b** with the rendering space **112** of FIG. **1A**. For example, the sensor signal processor **130** of FIG. **1B** may receive the sensor signals.

(80) The touch event may be associated with a contact space defined with respect to the device, based on the sensor signal (**204**). For example, the touch events **114**, **118** may be related to respective, corresponding contact spaces **116**, **120** of the contact map **110** of FIG. **1A**. For example, the gesture recognition model **132** may process outputs of the sensor signal preprocessor **130** to determine an image representation of the contact map **110**, including the contact spaces **116**, **120**.

(81) A user interface event of the user interface may be rendered within the rendering space that corresponds to the contact space (**206**). For example, the user interface **102b** may have the user interface events **116a**, **120a** rendered within the rendering space **112**. For example, the mapping module **134** of FIG. **1B** may map the contact spaces **116**, **120** at locations within the contact map **110** to respective, corresponding user interface events **116a**, **120a** within corresponding spatial locations within the rendering space **112**, where the rendering engine **136** of FIG. **1B** may be configured to execute the rendering of the user interface **102b** with respect to the rendering space **112**.

(82) FIG. 3 is a block diagram of an alternate example of the system of FIG. **1B**. In FIG. 3, an inertial measurement unit (IMU) **302** represents an example of one of the motion sensors **106**, **108** of FIG. **1A**. The IMU **302** may include or have access to, e.g., an accelerometer, gyroscope, magnetometer, or other motion-related signal information. For example, the IMU **302** may provide measurements for three degrees of freedom for each sensor component, including, e.g., measurements related to acceleration (in X, Y, Z coordinates); gyroscope (as pitch, yaw, roll), rotational acceleration, and various other features, each of which may be sampled with a designated sampling rate frequency.

(83) Thus, the IMU **302** is illustrated as outputting various sensor signals **304** over a period of time. In FIG. 3, in more detail, successive, overlapping time windows **306** of sensor signals **304** may be collected and buffered for storage within a feature map **308**. For example, the feature map **308** may represent an $m \times n$ matrix of 'm' features over 'n' time periods that effectively reformats the sensor

signal(s) **304** into a matrix (or vector) format.

(84) An infinite impulse response (IIR) filter **310** provides an example of the sensor signal preprocessor **130** of FIG. **1B**. As discussed with respect to the sensor signal preprocessor **130**, the IIR filter **310** may be configured to filter aspects or portions of the sensor signal(s) **304** that correspond to user movements not related to the gesture detection and implementation techniques described herein, or to calibrate operations for use with a particular implementation of the HMD **102**.

(85) The feature map **308** and the IIR filter **310** may also be configured to be compatible with a gesture detection model **312**. That is, as referenced above, the system of FIG. **3** may be configured to adapt an existing gesture detection model for use with the techniques of FIGS. **1A**, **1B**, and **2**, and for the particular user **104** and the particular HMD **102**. For example, the feature map **308** and the IIR filter **310** may express desired features or aspects of the sensor signals **304** in a format that is compatible with the gesture detection model **312**.

(86) In the example of FIG. **3**, the feature map **308** and the IIR filter **310** may primarily be provided for signal formatting and, e.g., noise reduction, where operations of the IIR filter **310** in providing noise reduction that is customized and calibrated, as just referenced. In addition, as also referenced, the IMU **302** and other motion sensors may have a mounting orientation(s) that varies among different HMD implementations. Additional or alternative signal pre-processing may be performed to, e.g., virtually re-mount an IMU(s) on different HMDs to provide reorientation towards a common reference frame.

(87) Such types of calibration, including calibrating the IIR filter **310** as well, may be performed in a factory setting to provide desired virtual re-mounting for compatibility with the gesture detection model **312** being used. These and related calibration techniques provide highly customizable implementations of the user interface controller **128**, without requiring training of a new/different implementation of the gesture detection model **134** of FIG. **1B** or the gesture detection model **312** of FIG. **3** for each new/different HMD.

(88) Some aspects of calibration may also be performed by the individual user **104**. For example, upon initial use of the HMD **102**, a noise floor for the sensor signals **304** may be initially obtained to obtain spectrum calibration data, and the IIR filter **310** may then be fit to the spectrum calibration data to obtain a desired signal to noise ratio.

(89) Thus, use of the IMU **302** enables fine-grained detection of characteristics of touch events. For example, the IMU **302**, being rigidly mounted to the HMD **102**, is able to capture multiple types of touch interactions, including taps, swipes, pinches, rotations, holds, and various other types of interactions. For example, a tap may be detected as a vibration at a particular angle, while finger movement in an orthogonal direction is registered along a different axis of the IMU **302**.

(90) FIG. **4** is a block diagram of an alternate example of the system of FIG. **1B**. In the example of FIG. **4**, as in FIG. **3**, IMU signals are frequently captured and sampled as new updates arrive. In FIG. **4**, an array of IMUs dispersed over a HMD, such as the HMD **102**, is used.

(91) In particular, as shown, FIG. **4** illustrates a first IMU **402a** and an Nth IMU **402b** of the array. The IMUs **402a**, **402b** are illustrated as outputting respective sets of sensor signals **404a**, **404b** over a period of time. As in FIG. **3**, successive, overlapping time windows **406a**, **406b** of sensor signals **404a**, **404b** may be collected and buffered for storage within corresponding feature maps **408a**, **408b**. IIR filters **410a**, **410b**, and potentially other types of sensor signal preprocessing as described above, may be provided for corresponding IMUs **402a**, **402b**.

(92) The feature maps **408a**, **408b** and the IIR filters **410a**, **410b** may also be configured to be compatible with a gesture detection model **312**. That is, as referenced above, the system of FIG. **3** may be configured to adapt an existing gesture detection model for use with the techniques of FIGS. **1A**, **1B**, and **2**, and for the particular user **104** and the particular HMD **102**. For example, the feature map **308** and the IIR filter **310** may express desired features or aspects of the sensor signals **304** in a format that is compatible with the gesture detection model **312**.

(93) In FIG. 4, a gesture detection model **414** includes a multi-head architecture in which an attention head **411a** and an attention head **411b** receive filtered outputs of the IIR filters **410a**, **410b**, respectively. For example, the attention heads **411a**, **411b** may independently calculate attention weights for their corresponding inputs, which may then be combined and processed through a fully connected layer(s) of an encoder **412**.

(94) The encoder **412** may be included with a corresponding decoder **418** in an autoencoder bottleneck **416** for use in mapping detected contact spaces to corresponding rendering spaces, as described above with respect to FIG. 1. In more detail, the autoencoder bottleneck **416** refers to a technique for turning a relatively high-dimensional input into a latent low-dimensional code (with encoder **412**) and then provide a reconstruction using the latent code at the decoder **418**.

(95) For example, the encoder **412** may be configured as a set of convolutional blocks and pooling modules to compress combined inputs from the attention heads **411a**, **411b**. The decoder **418** may include upsampling modules and convolutional blocks that restore or reconstruct the input (e.g., with reduced noise). By restricting a flow of information, the autoencoder bottleneck **416** effectively enables the most important, e.g., desired, information to pass through, thereby forming a knowledge representation of the input and providing correlations between different received inputs. Put another way, the decoder **418** decompresses a compressed knowledge representation to provide an improved and more useful image from latent attributes of the image.

(96) Thus, the multi-head architecture of FIG. 4 provides an image to image translator at an end stage to estimate contact spaces on a contact map, which may then be processed by map assignment logic to relate device surfaces of the HMD **102** to the estimated contact map, and thereby to a rendering space of a user interface of the HMD **102**.

(97) FIG. 5 illustrates an example use case scenario for the systems of FIGS. 1B and 4. In FIG. 5, a device surface **502** may represent, e.g., a face of a pair of smart goggles, which may be recognized as a contact map (similar to the contact map **110** of FIG. 1A). Then, a touch event **503** may be recognized through the use of a pool of IMU sensors **504a**, **504b**, **504c**, **504d** that are dispersed around a perimeter of the device surface **502**.

(98) For example, the touch event **503** may include placement of a single finger on the device surface. Such finger placement may be detected by each of the IMUs **504a**, **504b**, **504c**, **504d**, but with differing amplitudes that vary based on proximity of each IMU to the touch event **503**.

(99) For example, as shown, the IMU **504c** closest to the touch event **503** may experience the largest amplitude, indicated by a relative size of an arrow **506c**. The IMU **504a**, being next-closest to the touch event **503**, may experience a second largest amplitude, indicated by a relative size of an arrow **506a**. More distance IMUs **504b** and **504d** may detect signals with correspondingly smaller amplitudes, represented by relative sizes of respective arrows **506b** and **506d**.

(100) In the example of FIG. 5, a user interface **508** visible to the user **104** when using/wearing the smart goggles may be providing with a VR environment **510** as an example of the rendering space **112** of FIG. 1A. Then, a punch through area **512** may be provided in response to the touch event **503**, so that the user **104** may be provided with a window through the VR environment **510** and into a surrounding real-world environment.

(101) As described, the punch through area **512** may be provided with a size, and in a location, corresponding to the touch event **503**. For example, the touch event **503** occurring on a right side of the user **104** may result in the punch through area **512** also occurring on the right side of the user **104**. Of course, any desired mapping may be used to relate touch events with subsequent user interface events such as the punch through area **512**.

(102) In the example of FIG. 6, a touch event includes two fingers **603a**, **603b** being placed on the device surface (contact map) **502**. As in FIG. 5, such finger placement may be detected by each of the IMUs **504a**, **504b**, **504c**, **504d**, but with differing amplitudes that vary based on proximity of each IMU to the touch event of fingers **603a**, **603b**.

(103) For example, as shown, the closest IMU **504c** may experience the largest amplitude,

indicated by a relative size of an arrow **606c**, which may be larger than the amplitude indicated by the arrow **506c** of FIG. 5. Further in FIG. 6, amplitudes of the IMUs **504a**, **504d**, indicated by respective arrows **606a**, **606d**, are smaller than the amplitude indicated by the arrow **606c**, but are both larger than corresponding amplitudes/arrows **506a**, **506d** in FIG. 5. Similarly, the amplitude indicated by the arrow **606b** is the smallest of FIG. 6, but larger than that of corresponding arrow **506b** of FIG. 5.

(104) Thus, FIG. 6 illustrates that the distributed pool of IMU sensors **504a**, **504b**, **504c**, **504d** may be used by described techniques to determine fine-grained nuances and differences between touch events **503** and **603a/603b**, even when touch events are very similar in type or kind. Such differences may be usefully reflected in operation of the user interface **508**, e.g., of the VR environment **510**.

(105) For example, in FIG. 6, a resulting punch through area **612** may be larger than a punch through area **512** of FIG. 5. Moreover, additional or alternative finger placement/movement aspects may be used for desired effects. For example, placement of three fingers may result in yet a larger punch through area. In other examples, horizontal (as compared to vertical) finger placement may have a specified effect. Additionally, as described above, the device surface **502** should be understood to potentially represent two or more device surfaces, and each such device surface (at its respective location) may be assigned different mappings and gestures, as desired.

(106) Moreover, FIGS. 5 and 6 illustrate that finger movements or changes over both a space of the device surface (contact map) **502** and over a time during which the touch event(s) occurs may be used to provide spatiotemporally complex gestures. For example, the user **104** may exhibit finger movement across a face of the device surface **502** at fast or slow speeds, either linearly or rotationally, and combinations of two or more fingers may be used to indicate desired gestures for corresponding rendering effects.

(107) FIG. 7 is a third person view of a user **702** (analogous to the user **104** of FIG. 1) in an ambient environment **7000**, with one or more external computing systems shown as additional resources **752** that are accessible to the user **702** via a network **7200**. FIG. 7 illustrates numerous different wearable devices that are operable by the user **702** on one or more body parts of the user **702**, including a first wearable device **750** in the form of glasses worn on the head of the user, a second wearable device **754** in the form of ear buds worn in one or both ears of the user **702**, a third wearable device **756** in the form of a watch worn on the wrist of the user, and a computing device **706** held by the user **702**. In FIG. 7, the computing device **706** is illustrated as a handheld computing device but may also be understood to represent any personal computing device, such as a table or personal computer.

(108) In some examples, the first wearable device **750** is in the form of a pair of smart glasses including, for example, a display, one or more image sensors that can capture images of the ambient environment, audio input/output devices, user input capability, computing/processing capability and the like. Additional examples of the first wearable device **750** are provided below, with respect to FIGS. 8A and 8B.

(109) In some examples, the second wearable device **754** is in the form of an ear worn computing device such as headphones, or earbuds, that can include audio input/output capability, an image sensor that can capture images of the ambient environment **7000**, computing/processing capability, user input capability and the like. In some examples, the third wearable device **756** is in the form of a smart watch or smart band that includes, for example, a display, an image sensor that can capture images of the ambient environment, audio input/output capability, computing/processing capability, user input capability and the like. In some examples, the handheld computing device **706** can include a display, one or more image sensors that can capture images of the ambient environment, audio input/output capability, computing/processing capability, user input capability, and the like, such as in a smartphone. In some examples, the example wearable devices **750**, **754**, **756** and the example handheld computing device **706** can communicate with each other and/or with external

computing system(s) **752** to exchange information, to receive and transmit input and/or output, and the like. The principles to be described herein may be applied to other types of wearable devices not specifically shown in FIG. 7 or described herein.

(110) The user **702** may choose to use any one or more of the devices **706**, **750**, **754**, or **756**, perhaps in conjunction with the external resources **752**, to implement any of the implementations described above with respect to FIGS. 1-6. For example, the user **702** may use an application executing on the device **706** and/or the smartglasses **750** to execute the user interface controller **128** of FIG. 1B.

(111) The device **706** may access the additional resources **752** to facilitate the aspects of the various techniques described herein, or related techniques. In some examples, the additional resources **752** may be partially or completely available locally on the device **706**. In some examples, some of the additional resources **752** may be available locally on the device **706**, and some of the additional resources **752** may be available to the device **706** via the network **7200**. As shown, the additional resources **752** may include, for example, server computer systems, processors, databases, memory storage, and the like. In some examples, the processor(s) may include training engine(s), transcription engine(s), translation engine(s), rendering engine(s), and other such processors.

(112) The device **706** may operate under the control of a control system **760**. The device **706** can communicate with one or more external devices, either directly (via wired and/or wireless communication), or via the network **7200**. In some examples, the one or more external devices may include various ones of the illustrated wearable computing devices **750**, **754**, **756**, another mobile computing device similar to the device **706**, and the like. In some implementations, the device **706** includes a communication module **762** to facilitate external communication. In some implementations, the device **706** includes a sensing system **764** including various sensing system components. The sensing system components may include, for example, one or more image sensors **765**, one or more position/orientation sensor(s) **764** (including for example, an inertial measurement unit, an accelerometer, a gyroscope, a magnetometer and other such sensors), one or more audio sensors **766** that can detect audio input, one or more image sensors **767** that can detect visual input, one or more touch input sensors **768** that can detect touch inputs, and other such sensors. The device **706** can include more, or fewer, sensing devices and/or combinations of sensing devices. Various ones of the various sensors may be used individually or together to perform the types of saliency detection described herein.

(113) Captured still and/or moving images may be displayed by a display device of an output system **772**, and/or transmitted externally via a communication module **762** and the network **7200**, and/or stored in a memory **770** of the device **706**. The device **706** may include one or more processor(s) **774**. The processors **774** may include various modules or engines configured to perform various functions. In some examples, the processor(s) **774** may include, e.g., training engine(s), transcription engine(s), translation engine(s), rendering engine(s), and other such processors. The processor(s) **774** may be formed in a substrate configured to execute one or more machine executable instructions or pieces of software, firmware, or a combination thereof. The processor(s) **774** can be semiconductor-based including semiconductor material that can perform digital logic. The memory **770** may include any type of storage device or non-transitory computer-readable storage medium that stores information in a format that can be read and/or executed by the processor(s) **774**. The memory **770** may store applications and modules that, when executed by the processor(s) **774**, perform certain operations. In some examples, the applications and modules may be stored in an external storage device and loaded into the memory **770**.

(114) Although not shown separately in FIG. 7, it will be appreciated that the various resources of the computing device **706** may be implemented in whole or in part within one or more of various wearable devices, including the illustrated smartglasses **750**, earbuds **754**, and smartwatch **756**, which may be in communication with one another to provide the various features and functions described herein.

(115) An example head mounted wearable device **800** in the form of a pair of smart glasses is shown in FIGS. **8A** and **8B**, for purposes of discussion and illustration. The example head mounted wearable device **800** includes a frame **802** having rim portions **803** surrounding glass portion, or lenses **807**, and arm portions **830** coupled to a respective rim portion **803**. In some examples, the lenses **807** may be corrective/prescription lenses. In some examples, the lenses **807** may be glass portions that do not necessarily incorporate corrective/prescription parameters. A bridge portion **809** may connect the rim portions **803** of the frame **802**. In the example shown in FIGS. **8A** and **8B**, the wearable device **800** is in the form of a pair of smart glasses, or augmented reality glasses, simply for purposes of discussion and illustration.

(116) In some examples, the wearable device **800** includes a display device **804** that can output visual content, for example, at an output coupler providing a visual display area **805**, so that the visual content is visible to the user. In the example shown in FIGS. **8A** and **8B**, the display device **804** is provided in one of the two arm portions **830**, simply for purposes of discussion and illustration. Display devices **804** may be provided in each of the two arm portions **830** to provide for binocular output of content. In some examples, the display device **804** may be a see through near eye display. In some examples, the display device **804** may be configured to project light from a display source onto a portion of teleprompter glass functioning as a beamsplitter seated at an angle (e.g., 30-45 degrees). The beamsplitter may allow for reflection and transmission values that allow the light from the display source to be partially reflected while the remaining light is transmitted through. Such an optic design may allow a user to see both physical items in the world, for example, through the lenses **807**, next to content (for example, digital images, user interface elements, virtual content, and the like) output by the display device **804**. In some implementations, waveguide optics may be used to depict content on the display device **804**.

(117) The example wearable device **800**, in the form of smart glasses as shown in FIGS. **8A** and **8B**, includes one or more of an audio output device **806** (such as, for example, one or more speakers), an illumination device **808**, a sensing system **810**, a control system **812**, at least one processor **814**, and an outward facing image sensor **816** (for example, a camera). In some examples, the sensing system **810** may include various sensing devices and the control system **812** may include various control system devices including, for example, the at least one processor **814** operably coupled to the components of the control system **812**. In some examples, the control system **812** may include a communication module providing for communication and exchange of information between the wearable device **800** and other external devices. In some examples, the head mounted wearable device **800** includes a gaze tracking device **815** to detect and track eye gaze direction and movement. Data captured by the gaze tracking device **815** may be processed to detect and track gaze direction and movement as a user input. In the example shown in FIGS. **8A** and **8B**, the gaze tracking device **815** is provided in one of two arm portions **830**, simply for purposes of discussion and illustration. In the example arrangement shown in FIGS. **8A** and **8B**, the gaze tracking device **815** is provided in the same arm portion **830** as the display device **804**, so that user eye gaze can be tracked not only with respect to objects in the physical environment, but also with respect to the content output for display by the display device **804**. In some examples, gaze tracking devices **815** may be provided in each of the two arm portions **830** to provide for gaze tracking of each of the two eyes of the user. In some examples, display devices **804** may be provided in each of the two arm portions **830** to provide for binocular display of visual content.

(118) The wearable device **800** is illustrated as glasses, such as smartglasses, augmented reality (AR) glasses, or virtual reality (VR) glasses. More generally, the wearable device **800** may represent any head-mounted device (HMD), including, e.g., goggles, helmet, or headband, or another device (e.g., smartphone) mounted to a user's head using an external frame for use as a HMD. Even more generally, the wearable device **800** and the computing device **706** may represent any wearable device(s), handheld computing device(s), or combinations thereof.

(119) Use of the wearable device **800**, and similar wearable or handheld devices such as those

shown in FIG. 7, enables useful and convenient use case scenarios of implementations of FIGS. 1-6. For example, any specified portions of the frame **802** and/or the arm portions **830** may be designated as a contact map(s), such as the contact map **110** of FIG. 1. For example, a portion of the frame **802** enclosing the lenses **807** may be used to control a cursor or other control aspect of user interface in a horizontal plane (x, y coordinates), while one or both of the arm portions **830** may be used to control a cursor or other control aspect of the same user interface in a depth direction (z coordinates).

(120) Various implementations of the systems and techniques described here can be realized in digital electronic circuitry, integrated circuitry, specially designed ASICs (application specific integrated circuits), computer hardware, firmware, software, and/or combinations thereof. These various implementations can include implementation in one or more computer programs that are executable and/or interpretable on a programmable system including at least one programmable processor, which may be special or general purpose, coupled to receive data and instructions from, and to transmit data and instructions to, a storage system, at least one input device, and at least one output device.

(121) These computer programs (also known as modules, programs, software, software applications or code) include machine instructions for a programmable processor, and can be implemented in a high-level procedural and/or object-oriented programming language, and/or in assembly/machine language. As used herein, the terms “machine-readable medium” “computer-readable medium” refers to any computer program product, apparatus and/or device (e.g., magnetic discs, optical disks, memory, Programmable Logic Devices (PLDs)) used to provide machine instructions and/or data to a programmable processor, including a machine-readable medium that receives machine instructions as a machine-readable signal. The term “machine-readable signal” refers to any signal used to provide machine instructions and/or data to a programmable processor.

(122) To provide for interaction with a user, the systems and techniques described here can be implemented on a computer having a display device (e.g., a CRT (cathode ray tube) or LCD (liquid crystal display) monitor, or LED (light emitting diode)) for displaying information to the user and a keyboard and a pointing device (e.g., a mouse or a trackball) by which the user can provide input to the computer. Other kinds of devices can be used to provide for interaction with a user as well. For example, feedback provided to the user can be any form of sensory feedback (e.g., visual feedback, auditory feedback, or tactile feedback), and input from the user can be received in any form, including acoustic, speech, or tactile input.

(123) The systems and techniques described here can be implemented in a computing system that includes a back end component (e.g., as a data server), or that includes a middleware component (e.g., an application server), or that includes a front end component (e.g., a client computer having a graphical user interface or a Web browser through which a user can interact with an implementation of the systems and techniques described here), or any combination of such back end, middleware, or front end components. The components of the system can be interconnected by any form or medium of digital data communication (e.g., a communication network). Examples of communication networks include a local area network (“LAN”), a wide area network (“WAN”), and the Internet.

(124) The computing system can include clients and servers. A client and server are generally remote from each other and typically interact through a communication network. The relationship of client and server arises by virtue of computer programs running on the respective computers and having a client-server relationship to each other.

(125) A number of implementations have been described. Nevertheless, it will be understood that various modifications may be made without departing from the spirit and scope of the description and claims.

(126) In addition, the logic flows depicted in the figures do not require the particular order shown, or sequential order, to achieve desirable results. In addition, other steps may be provided, or steps

may be eliminated, from the described flows, and other components may be added to, or removed from, the described systems. Accordingly, other implementations are within the scope of the following claims.

(127) Further to the descriptions above, a user is provided with controls allowing the user to make an election as to both if and when systems, programs, devices, networks, or features described herein may enable collection of user information (e.g., information about a user's social network, social actions, or activities, profession, a user's preferences, or a user's current location), and if the user is sent content or communications from a server. In addition, certain data may be treated in one or more ways before it is stored or used, so that user information is removed. For example, a user's identity may be treated so that no user information can be determined for the user, or a user's geographic location may be generalized where location information is obtained (such as to a city, ZIP code, or state level), so that a particular location of a user cannot be determined. Thus, the user may have control over what information is collected about the user, how that information is used, and what information is provided to the user.

(128) The computer system (e.g., computing device) may be configured to wirelessly communicate with a network server over a network via a communication link established with the network server using any known wireless communications technologies and protocols including radio frequency (RF), microwave frequency (MWF), and/or infrared frequency (IRF) wireless communications technologies and protocols adapted for communication over the network.

(129) In accordance with aspects of the disclosure, implementations of various techniques described herein may be implemented in digital electronic circuitry, or in computer hardware, firmware, software, or in combinations of them. Implementations may be implemented as a computer program product (e.g., a computer program tangibly embodied in an information carrier, a machine-readable storage device, a computer-readable medium, a tangible computer-readable medium), for processing by, or to control the operation of, data processing apparatus (e.g., a programmable processor, a computer, or multiple computers). In some implementations, a tangible computer-readable storage medium may be configured to store instructions that when executed cause a processor to perform a process. A computer program, such as the computer program(s) described above, may be written in any form of programming language, including compiled or interpreted languages, and may be deployed in any form, including as a standalone program or as a module, component, subroutine, or other unit suitable for use in a computing environment. A computer program may be deployed to be processed on one computer or on multiple computers at one site or distributed across multiple sites and interconnected by a communication network.

(130) Specific structural and functional details disclosed herein are merely representative for purposes of describing example implementations. Example implementations, however, may be embodied in many alternate forms and should not be construed as limited to only the implementations set forth herein.

(131) The terminology used herein is for the purpose of describing particular implementations only and is not intended to be limiting of the implementations. As used herein, the singular forms “a,” “an,” and “the” are intended to include the plural forms as well, unless the context clearly indicates otherwise. It will be further understood that the terms “comprises,” “comprising,” “includes,” and/or “including,” when used in this specification, specify the presence of the stated features, steps, operations, elements, and/or components, but do not preclude the presence or addition of one or more other features, steps, operations, elements, components, and/or groups thereof.

(132) It will be understood that when an element is referred to as being “coupled,” “connected,” or “responsive” to, or “on,” another element, it can be directly coupled, connected, or responsive to, or on, the other element, or intervening elements may also be present. In contrast, when an element is referred to as being “directly coupled,” “directly connected,” or “directly responsive” to, or “directly on,” another element, there are no intervening elements present. As used herein the term “and/or” includes any and all combinations of one or more of the associated listed items.

(133) Spatially relative terms, such as “beneath,” “below,” “lower,” “above,” “upper,” and the like, may be used herein for ease of description to describe one element or feature in relationship to another element(s) or feature(s) as illustrated in the figures. It will be understood that the spatially relative terms are intended to encompass different orientations of the device in use or operation in addition to the orientation depicted in the figures. For example, if the device in the figures is turned over, elements described as “below” or “beneath” other elements or features would then be oriented “above” the other elements or features. Thus, the term “below” can encompass both an orientation of above and below. The device may be otherwise oriented (rotated 130 degrees or at other orientations) and the spatially relative descriptors used herein may be interpreted accordingly.

(134) Example implementations of the concepts are described herein with reference to cross-sectional illustrations that are schematic illustrations of idealized implementations (and intermediate structures) of example implementations. As such, variations from the shapes of the illustrations as a result, for example, of manufacturing techniques and/or tolerances, are to be expected. Thus, example implementations of the described concepts should not be construed as limited to the particular shapes of regions illustrated herein but are to include deviations in shapes that result, for example, from manufacturing. Accordingly, the regions illustrated in the figures are schematic in nature and their shapes are not intended to illustrate the actual shape of a region of a device and are not intended to limit the scope of example implementations.

(135) It will be understood that although the terms “first,” “second,” etc. may be used herein to describe various elements, these elements should not be limited by these terms. These terms are only used to distinguish one element from another. Thus, a “first” element could be termed a “second” element without departing from the teachings of the present implementations.

(136) Unless otherwise defined, the terms (including technical and scientific terms) used herein have the same meaning as commonly understood by one of ordinary skill in the art to which these concepts belong. It will be further understood that terms, such as those defined in commonly used dictionaries, should be interpreted as having a meaning that is consistent with their meaning in the context of the relevant art and/or the present specification and will not be interpreted in an idealized or overly formal sense unless expressly so defined herein.

(137) While certain features of the described implementations have been illustrated as described herein, many modifications, substitutions, changes, and equivalents will now occur to those skilled in the art. It is, therefore, to be understood that the appended claims are intended to cover such modifications and changes as fall within the scope of the implementations. It should be understood that they have been presented by way of example only, not limitation, and various changes in form and details may be made. Any portion of the apparatus and/or methods described herein may be combined in any combination, except mutually exclusive combinations. The implementations described herein can include various combinations and/or sub-combinations of the functions, components, and/or features of the different implementations described.

Claims

1. A method comprising: receiving, in response to a touch event, a sensor signal from a motion sensor coupled to a device, the device providing a user interface with a rendering space that is defined relative to the device; associating the touch event with a contact space defined with respect to the device, based on the sensor signal; and rendering a user interface event of the user interface within the rendering space that corresponds to the contact space.
2. The method of claim 1, further comprising: rendering the user interface event at a rendering location within the rendering space that corresponds to a contact space location within a contact map defined with respect to the device.
3. The method of claim 1, further comprising: receiving, in response to a second touch event, a second sensor signal; relating the second touch event to a second contact space defined with respect

to the device, based on the second sensor signal; and rendering a second user interface event of the user interface within the rendering space that corresponds to the second contact space.

4. The method of claim 3, further comprising: using the contact space to provide user interface control of the user interface, including the user interface event, in a horizontal plane; and using the second contact space to provide user interface control of the user interface, including the second user interface event, in a depth plane that is perpendicular to the horizontal plane.
5. The method of claim 1, further comprising: receiving at least two sensor signals from at least two motion sensors mounted on the device; and relating the touch event to the contact space, based on the at least two sensor signals.
6. The method of claim 1, further comprising: relating an area of the contact space to the user interface event.
7. The method of claim 1, further comprising: relating the touch event to the contact space using a gesture detection neural network.
8. The method of claim 1, wherein the device includes a head mounted device (HMD).
9. The method of claim 1, further comprising: receiving the sensor signal from an inertial measurement unit (IMU).
10. The method of claim 1, further comprising: defining a position of the contact space within a contact map defined with respect to the device.
11. The method of claim 1, further comprising: detecting a second touch event; and rendering a second user interface in response to the second touch event.
12. The method of claim 1, further comprising: detecting a swiping touch gesture in conjunction with the touch event; and rendering the user interface event including moving the user interface within the rendering space in correspondence with the swiping touch gesture.
13. The method of claim 1, further comprising: detecting a second contact space associated with a second touch event; and rendering the user interface event including rendering the user interface with a size determined from a distance between the contact space and the second contact space.
14. The method of claim 13, further comprising: detecting a change in the distance between the contact space and the second contact space; and changing the size of the user interface, based on the change in the distance.
15. The method of claim 1, wherein the user interface event includes closing the user interface.
16. The method of claim 1, wherein a size of a surface area of the contact space corresponding with a size of the user interface event within the rendering space.
17. A computer program product, the computer program product being tangibly embodied on a non-transitory computer-readable storage medium and comprising instructions that, when executed by at least one computing device, are configured to cause the at least one computing device to: receive, in response to a touch event, a sensor signal from a motion sensor coupled to a device, the device providing a user interface with a rendering space that is defined relative to the device; associate the touch event with a contact space defined with respect to the device, based on the sensor signal; and render a user interface event of the user interface within the rendering space that corresponds to the contact space.
18. The computer program product of claim 17, wherein the instructions, when executed by the at least one computing device, are further configured to cause the at least one computing device to: receive at least two sensor signals from at least two motion sensors mounted on the device; and relate the touch event to the contact space, based on the at least two sensor signals.
19. The computer program product of claim 17, wherein the instructions, when executed by the at least one computing device, are further configured to cause the at least one computing device to: relate an area of the contact space to the user interface event.
20. The computer program product of claim 17, wherein the instructions, when executed by the at least one computing device, are further configured to cause the at least one computing device to: relate the touch event to the contact space using a gesture detection neural network.

21. The computer program product of claim 17, wherein the instructions, when executed by the at least one computing device, are further configured to cause the at least one computing device to: render the user interface event at a rendering location within the rendering space that corresponds to a contact space location within a contact map defined with respect to the device.
22. The computer program product of claim 17, wherein the instructions, when executed by the at least one computing device, are further configured to cause the at least one computing device to: receive, in response to a second touch event, a second sensor signal; relate the second touch event to a second contact space defined with respect to the device, based on the second sensor signal; and render a second user interface event of the user interface within the rendering space that corresponds to the second contact space.
23. The computer program product of claim 22, wherein the instructions, when executed by the at least one computing device, are further configured to cause the at least one computing device to: use the contact space to provide user interface control of the user interface, including the user interface event, in a horizontal plane; and use the second contact space to provide user interface control of the user interface, including the second user interface event, in a depth plane that is perpendicular to the horizontal plane.
24. A wearable device comprising: at least one frame for positioning the wearable device on a body of a user; at least one display; at least one processor; and at least one memory, the at least one memory storing a set of instructions, which, when executed, cause the at least one processor to: receive, in response to a touch event, a sensor signal from a motion sensor coupled to the wearable device, the wearable device providing a user interface using the at least one display, with a rendering space that is defined relative to the wearable device; associate the touch event with a contact space defined with respect to the wearable device, based on the sensor signal; and render a user interface event of the user interface within the rendering space that corresponds to the contact space.
25. The wearable device of claim 24, wherein the set of instructions, when executed by the at least one processor, are further configured to cause the wearable device to: receive at least two sensor signals from at least two motion sensors mounted on the wearable device; and relate the touch event to the contact space, based on the at least two sensor signals.
26. The wearable device of claim 24, wherein the set of instructions, when executed by the at least one processor, are further configured to cause the wearable device to: render the user interface event at a rendering location within the rendering space that corresponds to a contact space location within a contact map defined with respect to a surface of the wearable device.
-