

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12399656
Kind Code	B2
Date of Patent	August 26, 2025
Inventor(s)	Ki; Yang Seok et al.

Network storage gateway

Abstract

A Solid State Drive (SSD) is disclosed. The SSD may include ports to receive requests from a host and to send requests to a second storage device. The SSD may include flash storage for data. An SSD controller may process the requests received from the host and generate the requests sent to the second storage device. The SSD may act as a cache for the second storage device.

Inventors:	Ki; Yang Seok (Palo Alto, CA), Kang; Yangwook (San Jose, CA)
Applicant:	Samsung Electronics Co., Ltd. (Suwon-si, KR)
Family ID:	1000008778234
Assignee:	SAMSUNG ELECTRONICS CO., LTD. (N/A, KR)
Appl. No.:	18/231241
Filed:	August 07, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20230384982 A1	Nov. 30, 2023

Related U.S. Application Data

continuation parent-doc US 17589886 20220131 US 11755254 child-doc US 18231241
continuation parent-doc US 16846256 20200410 US 11256448 20220222 child-doc US 17589886
us-provisional-application US 62948794 20191216

Publication Classification

Int. Cl.: G06F3/06 (20060101); G06F13/42 (20060101)

U.S. Cl.:

CPC **G06F3/0679** (20130101); **G06F3/0613** (20130101); **G06F3/0659** (20130101);
 G06F13/4282 (20130101); G06F2213/0026 (20130101)

Field of Classification Search

CPC: G06F (3/0679); G06F (3/0659); G06F (3/0613); G06F (2213/0026); G06F (2213/28);
 G06F (2213/2802); G06F (2213/2804)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7185062	12/2006	Lolayekar et al.	N/A	N/A
8601211	12/2012	Flynn et al.	N/A	N/A
8996781	12/2014	Schuetz et al.	N/A	N/A
9285995	12/2015	Malwankar	N/A	N/A
9430412	12/2015	Huang	N/A	N/A
10019362	12/2017	Chatterjee et al.	N/A	N/A
11256488	12/2021	Elango et al.	N/A	N/A
11755254	12/2022	Ki	710/308	G06F 13/387
2005/0201394	12/2004	Wang et al.	N/A	N/A
2007/0038833	12/2006	Yamamoto et al.	N/A	N/A
2008/0313364	12/2007	Flynn et al.	N/A	N/A
2009/0210620	12/2008	Jibbe et al.	N/A	N/A
2013/0086300	12/2012	Luca	N/A	N/A
2015/0067353	12/2014	Hui	N/A	N/A
2015/0160862	12/2014	Blott	711/148	G06F 3/0683
2015/0254003	12/2014	Lee et al.	N/A	N/A
2017/0091042	12/2016	Chou et al.	N/A	N/A
2017/0271030	12/2016	Li	N/A	N/A
2017/0286170	12/2016	De et al.	N/A	N/A
2017/0329640	12/2016	Bandic et al.	N/A	N/A
2018/0285019	12/2017	Olarig et al.	N/A	N/A
2018/0307650	12/2017	Kachare et al.	N/A	N/A
2019/0109720	12/2018	Olarig	N/A	N/A
2019/0141128	12/2018	Hallak et al.	N/A	N/A
2019/0146675	12/2018	Subramanian et al.	N/A	N/A
2019/0199690	12/2018	Klein	N/A	N/A
2019/0272245	12/2018	Olarig et al.	N/A	N/A
2019/0278507	12/2018	Du et al.	N/A	N/A
2020/0334292	12/2019	Ganeshan	N/A	G06F 3/0679

OTHER PUBLICATIONS

Notice of Allowance for U.S. Appl. No. 16/846,256, mailed Oct. 14, 2021. cited by applicant
Notice of Allowance for U.S. Appl. No. 17/589,886, mailed Apr. 27, 2023. cited by applicant
Office Action for U.S. Appl. No. 16/846,256, mailed Apr. 29, 2021. cited by applicant
Office Action for U.S. Appl. No. 17/589,886, mailed Oct. 14, 2022. cited by applicant

Primary Examiner: Phan; Raymond N

Attorney, Agent or Firm: Renaissance IP Law Group LLP

Background/Summary

RELATED APPLICATION DATA (1) This application is a continuation of U.S. patent application Ser. No. 17/589,886, filed Jan. 31, 2022, now allowed, which is a continuation of U.S. patent

application Ser. No. 16/846,256, filed Apr. 10, 2020, now U.S. Pat. No. 11,256,448, issued Feb. 22, 2022, which claims the benefit of U.S. Provisional Patent Application Ser. No. 62/948,794, filed Dec. 16, 2019, all of which are incorporated by reference herein for all purposes.

FIELD

(1) The inventive concepts relate generally to storage systems, and more particularly to storage systems that avoid bottlenecks in CPU and memory.

BACKGROUND

(2) Accessing data stored on storage devices in an efficient manner is becoming increasingly important. Storage devices are increasing in size, the amount of data stored on storage devices is also increasing, and therefore it becomes even more important to efficiently process and use this data.

(3) While historically storage has been handled separately from computing, with dedicated servers just for storage (and separate from computing servers), this arrangement has a potential bottleneck in the transmission of data between the computing servers and the storage servers. Thus, there has been a move to converge the computing server and the storage server into a single unit. But in a converged system, the central processing unit (CPU) becomes a bottleneck, which may slow down overall operations.

(4) A need remains to improve storage systems that avoid bottlenecks in CPU and memory.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1 shows a system including a compute server and a storage server, according to an embodiment of the inventive concept.

(2) FIG. 2 shows the conventional compute server of FIG. 1 using a converged architecture.

(3) FIGS. 3-5 shows the compute server and the storage server of FIG. 1 using diverged architectures.

(4) FIG. 6 shows details of the compute server and the storage server of FIG. 1.

(5) FIG. 7 shows messages being exchanged between the compute server of FIG. 1, the cache Solid State Drives (SSDs) of FIG. 5, and the Key-Value SSDs (KV-SSDs) of FIG. 5.

(6) FIG. 8 shows details of the cache SSDs of FIG. 5.

(7) FIG. 9 shows details of the SSDs of FIG. 5.

(8) FIG. 10 shows an alternative view of the cache SSDs of FIG. 5.

(9) FIGS. 11A-11C shows details of the cache namespace of FIG. 10.

(10) FIGS. 12A-12B show a flowchart of an example procedure for the cache SSD of FIG. 5 to process a write request from the compute server of FIG. 1, according to an embodiment of the inventive concept.

(11) FIGS. 13A-13B show a flowchart of an example procedure for the cache SSD of FIG. 5 to process a read request from the compute server of FIG. 1, according to an embodiment of the inventive concept.

DETAILED DESCRIPTION

(12) Reference will now be made in detail to embodiments of the inventive concept, examples of which are illustrated in the accompanying drawings. In the following detailed description, numerous specific details are set forth to enable a thorough understanding of the inventive concept. It should be understood, however, that persons having ordinary skill in the art may practice the inventive concept without these specific details. In other instances, well-known methods, procedures, components, circuits, and networks have not been described in detail so as not to unnecessarily obscure aspects of the embodiments.

(13) It will be understood that, although the terms first, second, etc. may be used herein to describe

various elements, these elements should not be limited by these terms. These terms are only used to distinguish one element from another. For example, a first module could be termed a second module, and, similarly, a second module could be termed a first module, without departing from the scope of the inventive concept.

(14) The terminology used in the description of the inventive concept herein is for the purpose of describing particular embodiments only and is not intended to be limiting of the inventive concept. As used in the description of the inventive concept and the appended claims, the singular forms “a”, “an”, and “the” are intended to include the plural forms as well, unless the context clearly indicates otherwise. It will also be understood that the term “and/or” as used herein refers to and encompasses any and all possible combinations of one or more of the associated listed items. It will be further understood that the terms “comprises” and/or “comprising,” when used in this specification, specify the presence of stated features, integers, steps, operations, elements, and/or components, but do not preclude the presence or addition of one or more other features, integers, steps, operations, elements, components, and/or groups thereof. The components and features of the drawings are not necessarily drawn to scale.

(15) The progression of technologies to be able to store and process data as data increases is expected to move through several variations, as shown in FIG. 1 below, with each variation solving some problems but potentially introducing new ones.

(16) A converged infrastructure puts everything—CPU, memory, cache, and storage devices all in the same box. This infrastructure has the advantage that with everything consolidated in a single box, it is easy to migrate both machine and data: there is no concern that, for example, the data might be left behind when the machine is migrated.

(17) But the converged infrastructure has two problems. First, the CPU is very busy. Aside from executing commands associated with the operating system and applications, the CPU must also execute software commands to manage accessing data from the storage devices.

(18) The second problem is that there is a limit to the number of SSDs that may be used in a converged infrastructure. A socket on the machine may only support up to 4 SSDs. Given that a typical machine only has two sockets, the machine itself may only support up to eight SSDs. As data grows, this number of SSDs may be insufficient for future needs.

(19) A solution to the problems of the converged infrastructure is to disaggregate computing and storage. In a disaggregated infrastructure, the storage devices are again housed in a storage server rather than being integrated into the box with the computing server. The computing server and the storage server may then communicate over a network, such as using Ethernet protocol.

(20) This infrastructure reduces the burden on the CPU, as it no longer has to handle communicating with the storage devices. This infrastructure also permits higher numbers of storage devices: 24, 36, or potentially more storage devices may be stored in a single storage server, thus overcoming the limit on the number of storage devices supported by the converged infrastructure.

(21) But the CPU is still busy handling the operating system and executing instructions for various applications, so the CPU burden is only reduced, it is not eliminated as a possible bottleneck. If too many applications are requesting services from the CPU, the CPU may still be a potential bottleneck. In addition, when Peripheral Component Interconnect Express (PCIe) Gen 5 becomes standard, storage devices using PCIe Gen 5 will be able to send and receive data at a rate of 400 GB/second: faster than memory can process data. Thus, the memory in the storage server may be a bottleneck.

(22) To address the potential problems of the disaggregated infrastructure, further modifications may be introduced. In a third architecture, the storage devices may be replaced with Ethernet SSDs, permitting remote access of the storage devices from the CPU. This change avoids the problem of memory becoming a bottleneck. But since this third architecture places responsibility for managing data read and writes back on the CPU, the CPU is once again burdened and may be a bottleneck, just as in the converged infrastructure. Further, because the storage devices are accessed remotely

over a network, this third architecture may introduce latency issues in communications between the CPU and the storage devices.

(23) Using a cache in this third architecture may both help and hinder operations. The cache may be able to provide data in a read request more quickly than a remote access to the Ethernet SSDs. But the cache must also be accessed when writing data, resulting in write requests that may make two or three hops to reach the SSD, rather than a more direct route, again impacting latency.

(24) To address the problems of this third architecture, an SSD with some compute processing capability may be used to the compute server in a fourth architecture. This SSD effectively becomes the cache of the compute server. This SSD, which may include a storage-related processing unit (SPU), may handle storage management, freeing the CPU of that burden. The SPU may also handle data-related commands, such as acceleration, compression, encryption, etc. Overhead may be further reduced by replacing block SSDs with key-value (KV) SSDs in the storage server. This fourth architecture avoids any problems arising from memory bottlenecks or latency, and reduces the burden on the CPU (although the burden is not entirely eliminated, since the CPU still executes instructions from applications). Other downsides to this fourth architecture include that the caches (if there are more than one) do not communicate with each other, and that the SPU is typically not as powerful as a CPU (and so may not necessarily be able to execute every instruction like a CPU, or may not be able to execute instructions as efficiently as a CPU).

(25) To address the problems with the fourth architecture, a fifth architecture may be used. In this fifth architecture the cache SSDs are enhanced with general data processing units (DPUs) that may operate similarly to the CPU. With the caches including DPUs, some or all of the applications may be shifted to the DPUs on the cache SSDs, freeing the host CPU to handle other tasks and eliminating the potential bottleneck of the CPU. In the fifth architecture (and the fourth architecture, as well), the SPUs may support block access to data and permit Erasure Coding or other techniques (such as Redundant Arrays of Independent Disks (RAIDs)) to spread data across multiple KV SSDs in the storage server.

(26) In addition, the KV SSDs may include SPU units themselves, to carry out acceleration, compression, encryption, etc., just like the SPUs of the cache SSDs in the fourth architecture. SPUs may also implement cache management features.

(27) SPUs may be implemented as part of a controller, using a Field Programmable Gate Array (FPGA), Application-Specific Integrated Circuit (ASIC), a Graphics Processing Unit (GPU), or any other desired implementation. DPUs, on the other hand, tend to need more processing power (since they may be asked to execute general instructions for an application), and therefore may be implemented using conventional processors.

(28) In the fourth and fifth architectures, the cache SSDs use both PCIe and Ethernet: PCIe for communicating with the host processor and Ethernet for communicating with the KV SSDs in the storage server. Thus, the cache SSDs may include two ports: one for PCIe communication and one for Ethernet communication. The cache SSDs may communicate with any number of KV SSDs via the Ethernet port: there is no requirement that a single cache SSD may only communicate with a single KV SSD (which would prevent the cache SSD SPU from offering Erasure Coding/RAID functionality).

(29) The cache SSDs may be organized using cache namespaces, with each namespace defining a logically separate area. Each namespace may have its own input and output queues. Since the cache SSDs use both PCIe and Ethernet, the cache namespaces have at least two queues: one on the PCIe side and one on the Ethernet side (there may be more than two queues, depending on how queues are structured: for example, each side has a submission queue and a completion queue: the submission queue is used to receive requests, and a completion queue is used to return results).

(30) The SPU may execute various command, such as Nop (do nothing), Insert (insert partial data into the PU buffer), Write Once (to reconstruct data from Erasure Coding), XOR, Extended OP (to do a more complicated data reconstruction), or Compression, to identify a few.

(31) FIG. 1 shows a system including a compute server and a storage server, according to an embodiment of the inventive concept. In FIG. 1, compute server **105** may be reachable from client **110** via network **115**. Client **110** and compute server **105** may be in a client-server relationship: client **110** may issue commands and compute server **105** may execute those commands. Alternatively, compute server **105** may be a computer being used directly by an end user, avoiding the involvement of client **110**. Compute server **105** may access storage server **120** for data via network **115**.

(32) Network **115** may be any variety or varieties of network. For example, network **115** may include a Local Area Network (LAN), a Wide Area Network (WAN), a Metropolitan Area Network (MAN), or a global network such as the Internet, among other possibilities. Data may be sent across network **115** directly, or it may be protected: for example, using encryption or a Virtual Private Network (VPN). Network **115** may include wired or wireless connections. In addition, network **115** may include any desired combinations of these alternatives. For example, client **110** might be connected via a wireless connection to a LAN that in turn connects via a wired connection to the Internet, which in turn connects to another LAN to which compute server **105** and storage server **120** are connected. The connections between compute server **105**, client **110**, and storage server **120** may vary: the connections do not have to be the same in all situations. Further, storage server **120** may be omitted from the diagram if compute server **105** includes the necessary storage.

(33) FIG. 2 shows the compute server **105** of FIG. 1 using a converged architecture. In FIG. 2, compute server **105** is shown as including processor **205**, memory **210**, cache **215**, and Peripheral Component Interconnect Express (PCIe) Solid State Drives (SSDs) **220-1** and **220-2**. Processor **205** may include a software stack, including the operating system, applications, storage software (such as a file system), and controller software to manage devices attached to compute server **105** (such as memory **210** and PCIe SSDs **220-1** and **220-2**). Processor **205** may be any variety of processor: for example, an Intel Xeon, Celeron, Itanium, or Atom processor, an AMD Opteron processor, an ARM processor, etc. While FIG. 2 shows a single processor **205**, compute server **105** may include any number of processors, each of which may be single core or multi-core processors, and may be mixed in any desired combination.

(34) Memory **210** and cache **215** may be conventional memory and cache used in support of processing on processor **205**. Memory **210** may be any variety of memory, such as flash memory, Dynamic Random Access Memory (DRAM), Static Random Access Memory (SRAM), Persistent Random Access Memory, Ferroelectric Random Access Memory (FRAM), or Non-Volatile Random Access Memory (NVRAM), such as Magnetoresistive Random Access Memory (MRAM) etc. Memory **210** may also be any desired combination of different memory types. Memory **210** may be managed by a memory controller (not shown in FIG. 2), which may be a separate component in compute server **105** with a driver that is part of the software stack. Memory **210** may be used to store data that may be termed “short-term”: that is, data not expected to be stored for extended periods of time. Examples of short-term data may include temporary files, data being used locally by applications (which may have been copied from other storage locations), and the like. Cache **215** may act as very fast local storage (potentially even faster than memory **210**) for data being used by processor **205**, and may include non-volatile storage.

(35) Processor **205** and memory **210** may also support an operating system under which various applications may be running. These applications may issue requests to read data from or write data to either memory **210** or PCIe SSDs **220-1** and **220-2**. Whereas memory **210** may be used to store data that may be termed “short-term”, PCIe SSDs **220-1** and **220-2** are storage devices that may be used to store data that is considered “long-term”: that is, data expected to be stored for extended periods of time. PCIe SSDs **220-1** and **220-2** may be accessed using controller software in the software stack running on processor **205**. While FIG. 2 shows two PCIe SSDs **220-1** and **220-2**, embodiments of the inventive concept may include storage devices of any type and connecting via any desired connection. Thus, PCIe SSDs **220-1** and **220-2** may be replaced with Serial AT

Attachment (SATA) hard disk drives, Ethernet SSDs, or storage devices of any other types. Further, embodiments of the inventive concept may include any number (zero or more) of storage devices, and each storage device may be of any desired type: thus, multiple different types of storage devices may be mixed in compute server **105**.

(36) Implementing compute server **105** as shown in FIG. 2 has some advantages. The design is simple (only one “box” to consider), and compute server **105** may be relocated with all data being migrated at the same time. But compute server **105** as shown in FIG. 2 also has some disadvantages. First, processor **205** is very busy, since it is responsible for managing all operations within compute server **105**: executing applications, running the operating system (and all its background operations), storing and retrieving data from PCIe SSDs **220-1** and **220-2**, and so on (this business may be termed CPU saturation). Further, compute server **105** typically only has two sockets, each of which usually supports no more than four storage devices. Thus, the total number of storage devices that may be installed in compute server **105** is typically eight, limiting the extent to which storage may be expanded within compute server **105**.

(37) To address some of the disadvantages of compute server **105** as shown in FIG. 2, the arrangement shown in FIG. 3 may be used. In FIG. 3, PCIe SSDs **220-1** and **220-2** of FIG. 2 have been removed from compute server **105**, and storage server **120** has been added. Storage server **120** may include processor **305**, memory **310**, cache **315**, and PCIe SSDs **320-1** and **320-2**. All of these components are functionally similar or identical to those described above with reference to FIG. 2, and their descriptions will not be repeated here. Further, as with FIG. 2, while FIG. 3 shows storage server **120** as including PCIe SSDs **320-1** and **320-2**, embodiments of the inventive concept may include any number of storage devices, of any types, and of mixed types. (In the remainder of this document, any reference to a particular type of storage device is intended to encompass any alternative types of storage devices, subject to any requirements of the described storage device.)

(38) By moving PCIe SSDs **220-1** and **220-2** of FIG. 2 from compute server **105** of FIG. 2 to PCIe SSDs **320-1** and **320-2** in storage server **120** in FIG. 3, some of the burdens on processor **205** of compute server **105** are reduced. The controller software may be moved from the software stack of processor **205** to processor **305**, reducing the load on processor **205**. Further, because storage server **120** may be designed for that purpose (rather than as a more general compute server), storage server **120** may support more than eight storage devices. For example, rack versions of storage devices may support 32 SSDs using M.2 connectors (even more than 32 devices may be possible), providing potentially significantly more storage expandability than may be offered by compute server **105** of FIG. 2.

(39) While FIG. 3 shows only one storage server **120**, embodiments of the inventive concept may include any number of storage servers. Further, compute server **105** may communicate with any and/or all of the available storage servers **120**, each of which may also communicate with each other.

(40) It is important to note that in FIG. 3, compute server **105** and storage server **120** communicate machine to machine. That is, compute server **105** and storage server **120** each have network ports, and data exchanged between the two machines may pass through the network ports via processors **205** and **305** and memories **210** and **310**. Thus, while FIG. 3 reduces the demand on processor **205** of compute server **105** of FIG. 2, the demand is not entirely eliminated: processor **205** may not have to manage the controller software in the software stack, but processor **205** still deals with the data flow. (Further, if there are multiple storage servers **120**, then the various storage servers **120** may communicate machine to machine as well.)

(41) In addition, looking to the future, when PCIe generation 5 becomes established, PCIe SSDs **320-1** and **320-2** may be able to send and receive data faster than memory **310** may be able to process that data. Thus, memory **310** may become a bottleneck for data delivery in the embodiment of the inventive concept shown in FIG. 4.

(42) To address some of the disadvantages of compute server **105** and storage server **120** as shown

in FIG. 3, the arrangement shown in FIG. 4 may be used. In FIG. 4, storage server 120 has been simplified, eliminating processor 305, memory 310, and cache 315. Instead, storage server 120 may include backplane 405, to which Ethernet SSDs 410-1 and 410-2 may be connected. Unlike storage server 120 of FIG. 3, where data passes through memory 310 of FIG. 3, Ethernet SSDs 410-1 and 410-2 permit direct communication to network 115 of FIG. 1. (Note that Ethernet SSDs 410-1 and 410-2 do necessarily have their own Ethernet network ports: storage server 120 may still include a network port. But Ethernet SSDs 410-1 and 410-2 are capable of communicating directly with the network port on storage server 120, without having to involve a processor and/or memory.) By eliminating memory 310 of FIG. 3, storage server 120 will not have a bottleneck introduced when PCIe Gen 5 is used.

(43) While FIG. 4 shows embodiments of the inventive concept as including Ethernet SSDs 410-1 and 410-2, embodiments of the inventive concept are not limited to Ethernet SSDs. Provided the storage devices include some protocol for communication with compute server 105 (and with the other network-connected storage devices), Ethernet SSDs 410-1 and 410-2 may be replaced with functionally equivalent storage devices.

(44) As with FIG. 3, there may be multiple storage servers 120 in FIG. 4. If so, then compute server 105 may communicate with any Ethernet SSDs 410-1 and 410-2 installed in any and/or all of storage servers 120. Further, Ethernet SSDs 410-1 and 410-2 may communicate with each other, and with Ethernet SSDs installed in other storage servers.

(45) But on the side of compute server 105, network traffic may still pass through processor 205 and memory 210. Thus, memory 210 may still be a potential bottleneck, with data being requested by or sent from Ethernet SSDs 410-1 and 410-2 faster than memory 210 may be able to process. In addition, with the controller software back in the software stack on processor 205, the burden on processor 205 may be increased (undoing some of the benefits of the embodiment of the inventive concept shown in FIG. 3). Finally, since compute server 105 and storage server 120 may communicate across a network, Input/Output latency (the time needed to complete communication) may become a factor in delaying data transmission.

(46) Thus, in FIG. 5, embodiments of the inventive concept are shown that may resolve these issues. In FIG. 5, cache 215 of FIG. 2 has been removed from compute server 105, and cache SSDs 505-1 and 505-2 have been added. Cache SSDs 505-1 and 505-2 may be SSDs (or other storage devices) that may act as local caches of data retrieved from storage server 120. (But compute server 105 may still include a local cache other than cache SSDs 505-1 and 505-2.) Cache SSDs 505-1 and 505-2 may be any desired variety of SSDs: for example, cache SSDs 505-1 and 505-2 may be PCIe SSDs (that is, SSDs that communicate with processor 205 over a PCIe bus). Cache SSDs 505-1 and 505-2 may therefore be block-based SSDs (although cache SSDs 505-1 and 505-2 may also use different storage methodologies than block-based storage). Cache SSDs 505-1 and 505-2 may be SSDs that are “smart”, in the sense that cache SSDs 505-1 and 505-2 may manage the storage and retrieval of data from storage server 120. To that end, cache SSDs 505-1 and 505-2 may include storage processing units (SPUs) 510-1 and 510-2. Cache SSDs 505-1 and 505-2 may also communicate directly over a network (such as network 115 of FIG. 1), much like Ethernet SSDs 410-1 and 410-2 of FIG. 4, rather than leaving communication to processor 205.

(47) SPUs 510-1 and 510-2 may handle any aspects of storage and retrieval of data. For example, SPUs 510-1 and 510-2 may handle encryption and decryption of data, compression and decompression of data, and the use of erasure coding and/or error correction codes (ECCs) when storing data on any and/or all storage devices within storage server 120. Thus, if the data is stored across multiple storage devices, SPUs 510-1 and 510-2 may select the storage devices on which the data is stored, how the data is divided across those storage devices (including the generation of any ECC data that may be stored on the storage devices), how the data is compressed, and how the data is encrypted. By including SPUs 510-1 and 510-2 in cache SSDs 505-1 and 505-2, the software stack within processor 205 is reduced (SPUs 510-1 and 510-2 act in place of the controller software

in terms of directing data to the appropriate storage devices in storage server **120**). SPUs may be implemented as part of a controller (such as a SSD controller within cache SSDs **505-1** and **505-2**), using a Field Programmable Gate Array (FPGA), Application-Specific Integrated Circuit (ASIC), a Graphics Processing Unit (GPU), a general purpose GPU (GPGPU), or any other desired implementation.

(48) Cache SSDs **505-1** and **505-2** may also include general data processing units (DPUs) **515-1** and **515-2**. While SPUs **510-1** and **510-2** enable the offloading of management of the storage devices within storage server **120**, DPUs **515-1** and **515-2** enable the offloading of application processing from processor **205**. That is, DPUs may be thought of as auxiliary processors that may be utilized by processor **205** to reduce its workload. Thus, DPUs **515-1** and **515-2** may be used to further reduce the burden on processor **205**, which may now be reduced to performing the basic management of compute server **105**.

(49) DPUs **515-1** and **515-2** may also be used to perform acceleration functions on data. For example, an application running on processor **205** might wish to perform a query on a database to count the number of entries that satisfy a particular criteria. The entirety of the data may be retrieved and loaded into memory **210** so that processor **205** may perform the query. But if the data is to be discarded immediately after the query is completed, such an approach requires transmission of a great deal of data when only a small amount of data is ultimately needed. Instead of having processor **205** perform the query, DPU **515-1** and/or **515-2** may perform the query on the data retrieved from storage server **120**, and just return to processor **205** the results of the query. (It may even be possible to shift performance of the query all the way to storage server **120** to further increase efficiencies, as discussed in co-pending U.S. patent application Ser. No. 16/820,665, filed Mar. 16, 2020, and co-pending U.S. patent application Ser. No. 16/820,675, filed Mar. 16, 2020, both of which claim the benefit of U.S. Provisional Patent Application Ser. No. 62/819,499, filed Mar. 15, 2019, U.S. Provisional Patent Application Ser. No. 62/834,900, filed Apr. 16, 2019, U.S. Provisional Patent Application Ser. No. 62/945,877, filed Dec. 9, 2019, and U.S. Provisional Patent Application Ser. No. 62/945,883, filed Dec. 9, 2019, all of which are incorporated by reference herein for all purposes.)

(50) While the above describes using DPUs **515-1** and **515-2** to perform the acceleration function, SPUs **510-1** and **510-2** may also be used to perform acceleration functions. But if SPUs **510-1** and **510-2** are not implemented using general processors, SPUs **510-1** and **510-2** may only be capable of implementing specific acceleration functions that for which they have been programmed. A general processor, such as may be used to implement DPUs **515-1** and **515-2**, may be able to perform any acceleration function without being specifically programmed to implement such acceleration functions.

(51) DPUs **515-1** and **515-2** are optional equipment: any number of cache SSDs (from zero to every cache SSD) in compute server **105** may include DPUs. Because DPUs **515-1** and **515-2** tend to need more processing power as compared with SPUs **510-1** and **510-2** (since they may be asked to execute general instructions for an application), DPUs **515-1** and **515-2** may be implemented using conventional processors included within cache SSDs **505-1** and **505-2**. In addition, cache SSDs **505-1** and **505-2** may include a single component that acts as both SPU and DPU, rather than separate components. Thus, for example, cache SSD **505-1** might include a single component that functions as both SPU **510-1** and DPU **515-1** (with cache SSD **505-2** potentially similarly arranged). As with other descriptions relating to storage devices, different cache SSDs may be implemented differently. Thus, one cache SSD in compute server **105** might include only an SPU, another cache SSD in compute server **105** might include both an SPU and a DPU, but implemented as separate components, and a third cache SSD in compute server **105** might include a single components that functions as both SPU and DPU.

(52) In FIG. 5, storage server **120** is shown as including Key-Value SSDs (KV-SSDs) **520-1** and **520-2**. While Ethernet SSDs **410-1** and **410-2** of FIG. 4 may be used in storage server **120** in the

embodiments of the inventive concept shown in FIG. 5, KV-SSDs tend to be more efficient than SSDs that use block-based storage. Thus, replacing Ethernet SSDs **410-1** and **410-2** of FIG. 4 with KV-SSDs **520-1** and **520-2** may introduce added efficiencies and reducing the workload on cache SSDs **505-1** and **505-2** (cache SSDs **505-1** and **505-2** do not necessarily need to track what is stored on which storage devices within storage server **120**) (although block-based storage may be used in the embodiments of the inventive concept shown in FIG. 5). Like Ethernet SSDs **410-1** and **410-2** of FIG. 4, KV-SSDs **520-1** and **520-2** may communicate directly with components of compute server **105** (and with other KV-SSDs, both within storage server **120** and other storage servers not shown in FIG. 5). Like cache SSDs **505-1** and **505-2**, KV-SSDs **520-1** and **520-2** may be any desired variety of SSD: for example, KV-SSDs **520-1** and **520-2** may be Ethernet KV-SSDs (that is, SSDs that are designed to communicate across networks using Ethernet).

(53) Like cache SSDs **505-1** and **505-2**, KV-SSDs **520-1** and **520-2** may include SPUs **525-1** and **525-2**. SPUs **525-1** and **525-2** may permit KV-SSDs **520-1** and **520-2** to perform local encryption/decryption, compression/decompression, error correction coding, acceleration, and other desired functions applicable to data stored on KV-SSDs **520-1** and **520-2**. SPUs **525-1** and **525-2** are optional equipment: any number of KV-SSDs (from zero to every KV-SSD) in storage server **120** may include SPUs. SPUs **525-1** and **525-2** may also be independent of SPUs **510-1** and **510-2**: any operations performed by SPUs **525-1** and **525-1** may be performed without reference to whether SPUs **510-1** and **510-2** have acted on the data. For example, co-pending U.S. patent application Ser. No. 16/820,665, filed Mar. 16, 2020, and co-pending U.S. patent application Ser. No. 16/820,675, filed Mar. 16, 2020, both of which claim the benefit of U.S. Provisional Patent Application Ser. No. 62/819,499, filed Mar. 15, 2019, U.S. Provisional Patent Application Ser. No. 62/834,900, filed Apr. 16, 2019, U.S. Provisional Patent Application Ser. No. 62/945,877, filed Dec. 9, 2019, and U.S. Provisional Patent Application Ser. No. 62/945,883, filed Dec. 9, 2019, all of which are incorporated by reference herein for all purposes, describe how data may be encoded as well as compressed, even though compression after encoding may not introduce much added benefit (in terms of space saved). In the same way, SPUs **525-1** and **525-2** may apply their own operations even though such operations may be of little benefit after the operations applied by SPUs **510-1** and **510-2**.

(54) FIG. 6 shows details of compute server **105** and storage server **120** of FIG. 1. In FIG. 6, typically, compute server **105** and storage server **120** may include one or more processors **205** and/or **305**, which may include memory controllers **605** and clocks **610**, which may be used to coordinate the operations of the components of the machine. Processors **205** and **305** may also be coupled to memories **210** and/or **310**, which may include random access memory (RAM), read-only memory (ROM), or other state preserving media, as examples. Processors **205** and/or **305** may also be coupled to storage devices **220**, **320**, **410**, **505**, and/or **520**, and to network connector **615**, which may be, for example, an Ethernet connector or a wireless connector. Processors **205** and/or **305** may also be connected to buses **620**, to which may be attached user interfaces **625** and Input/Output interface ports that may be managed using Input/Output engines **630**, among other components.

(55) FIG. 7 shows messages being exchanged between compute server **105** of FIG. 1, cache SSDs **505-1** and **505-2** of FIG. 5, and KV-SSDs **520-1** and **520-2** of FIG. 5. In FIG. 5, processor **205** of compute server **105** of FIG. 1 may send read request **705** and/or write request **710**. Requests **705** and **710** may be sent to cache SSD **505-1** (cache SSD **505-2** of FIG. 5 may also be the target of requests **705** and **710**, as may be any other cache SSDs in compute server **105** of FIG. 1). Cache SSD **505-1** may then receive requests **705** and **710** and attempt to satisfy them locally. For example, if cache SSD **505-1** stores data requested in read request **705**, cache SSD **505-1** may return that data directly to processor **205** as result **715**. Similarly, if cache SSD **505-1** receives write request **710**, cache SSD **505-1** may store the data in cache SSD **505-1** before sending any requests to KV-SSD **520-1** (or any other KV-SSD in any storage server **120** of FIG. 1). If so equipped, cache

SSD **505-1** may even send result **715** back to processor **205** before the data is ultimately written to KV-SSD **620-1** (which is sometimes termed an early acknowledgment).

(56) If cache SSD **505-1** does not store the data requested in read request **705**, or if cache SSD **505-1** receives write request **710**, cache SSD **505-1** may send read request **720** or write request **725** to KV-SSD **520-1**. KV-SSD **520-1** may then process requests **720** and **725** and may return result **730** (which may be the data requested in read request **720** or an acknowledgment that the data was written from write request **725**). Cache SSD **505-1** may then process result **730** and return result **715** to processor **205** in turn. (The above description refers to “processing” result **730** because if cache SSD **505-1** includes DPU **515-1** of FIG. 5, or if SPU **510-1** of FIG. 5 is properly implemented, cache SSD **505-1** may, for example, execute an acceleration function on the data returned in result **730**.)

(57) FIG. 8 shows details of cache SSDs **505-1** and **505-2** of FIG. 5. In FIG. 8, cache SSD **505-1** is shown. Cache SSD **505-1** may include ports **805** and **810**, SPU **510-1**, DPU **515-1**, persistent memory **815**, volatile memory **820**, and flash storage **825**. SPU **510-1** and DPU **515-1** have been discussed above with reference to FIG. 5 (recall that DPU **515-1** is optional, and may be omitted from cache SSD **505-1**).

(58) Ports **805** and **810** may be ports to communicate with various other components. For example, in one embodiment of the inventive concept, as discussed above, cache **505-1** may be a PCIe SSD and KV-SSDs **520-1** and **520-2** of FIG. 5 may be Ethernet-equipped KV-SSDs. Thus, port **805** may be a PCIe port for communicating with processor **205** of FIG. 5 and port **810** may be an Ethernet port for communicating with KV-SSDs **520-1** and **520-2** of FIG. 5.

(59) Persistent memory **815** may be used to store information that should not be lost, even if power were to be unexpectedly interrupted to cache SSD **505-1**. Examples of the information that might be maintained in persistent memory **815** may include information about which KV-SSDs **520-1** and **520-2** of FIG. 5 store data for which files (and where that data is stored on KV-SSDs **520-1** and/or **520-2**), or data that was subject to write request **710** of FIG. 7 from processor **205** of FIG. 5 but which has not yet been written to KV-SSDs **520-1** and/or **520-2** of FIG. 5. Because persistent memory **815** should be persistent, persistent memory **815** may be implemented using any desired technology that protects against data loss: for example, non-volatile storage or volatile storage that is subject to battery backup.

(60) Flash storage **825** may be used to store data as the cache for processor **205** of FIG. 5. As FIG. 8 describes cache SSD **505-1** as an SSD, and SSDs typically use flash storage for data, FIG. 8 shows cache SSD **505-1** as including flash storage **825**. But if cache SSD **505-1** were instead implemented using a different form of storage technology (for example, a hard disk drive), flash storage **825** may be replaced with the hardware appropriate to store data for that technology.

(61) Persistent memory **815** and flash storage **825** may seem redundant: typically both are forms of storage that are protected against data loss due to power interruptions. But persistent memory **815** be a faster form of storage than flash storage **825**: persistent memory **815** may be, for example, a cache for cache SSD **505-1**. But in some embodiments of the inventive concept, as flash storage **825** may be used to store all data in non-volatile forms, persistent memory **815** may be omitted.

(62) Volatile memory **820** may be used to store data whose loss due to an unexpected power interruption would not be a problem. For example, volatile memory **820** may be used to process data being subject to an acceleration function: if the data is only partially processed when power is lost, there may be little point in ensuring the data is not lost (it may not be possible to determine the portion of processing that has been completed on the data, and therefore the original data should be retrieved from KV-SSDs **520-1** and/or **520-2** again and the acceleration function executed from the start). Volatile memory **820** may also be used to store other data whose loss is not of concern.

Volatile memory **820** may be any form of memory as discussed above with reference to memories **210** and **310** of FIGS. 2-5.

(63) FIG. 9 shows details of SSDs **505-1**, **505-2**, **520-1**, and **520-2** of FIG. 5. In FIG. 9, SSDs **505-**

1, 505-2, 520-1, and 520-2 may include host interface layer (HIL) 905, SSD controller 910, and various flash memory chips 825-1 through 825-8 (also termed “flash memory storage”), which may be organized into various channels 915-1 through 915-4. Host interface logic 905 may manage communications between SSDs 505-1, 505-2, 520-1, and 520-2 and other components (such as processor 205 of FIG. 5 or other SSDs). These communications may include read requests 705 and/or 720 of FIG. 7 to read data from SSDs 505-1, 505-2, 520-1, and 520-2 and write requests 710 and/or 725 of FIG. 7 to write data to SSDs 505-1, 505-2, 520-1, and 520-2. Host interface logic 905 may manage an interface across only a single port (for example, KV-SSDs 520-1 and 520-2 of FIG. 5 might only communicate with cache SSDs 505-1 and 505-2 of FIG. 5), or it may manage interfaces across multiple ports (for example cache SSDs 505-1 and 505-2 of FIG. 5 may communicate using different approaches with both processor 205 of FIG. 5 and KV-SSDs 520-1 and 520-2 of FIG. 5). Alternatively, an SSD that includes multiple ports may have a separate host interface logic 905 to manage interfaces across each port. Embodiments of the inventive concept may also mix the possibilities (for example, an SSD with three ports might have one host interface logic to manage one port and a second host interface logic to manage the other two ports).

(64) SSD controller 910 may manage the read and write operations, along with garbage collection and other operations, on flash memory chips 825-1 through 825-8 using a flash memory controller (not shown in FIG. 9).

(65) SSD controller 910 may include mapping 920. Mapping 920 may manage where particular data is stored in flash memory chips 825-1 through 825-8. The exact form mapping 920 may take may depend on the implementation of the SSD. An SSD that uses block addressing may implement mapping 920 as a translation layer (sometimes referred to as flash translation layer (FTL)). A translation layer may perform the conventional functions of translating logical block addresses (LBAs) into physical block addresses (PBAs) where the data is actually stored. On the other hand, if SSD uses a key-value storage model, then mapping 920 may map object keys to addresses into physical addresses.

(66) While FIG. 9 shows SSDs 505-1, 505-2, 520-1, and 520-2 as including eight flash memory chips 825-1 through 825-8 organized into four channels 915-1 through 915-4, embodiments of the inventive concept may support any number of flash memory chips organized into any number of channels. Similarly, while FIG. 9 shows the structure of an SSD, other storage devices (for example, hard disk drives) may be implemented using a different structure.

(67) FIG. 10 shows an alternative view of cache SSDs 505-1 and 505-2 of FIG. 5. FIG. 10 presents more of an architectural/software view of cache SSDs 505-1 and 505-2: the hardware that supports this view is described further with reference to FIGS. 11A-11C below.

(68) In FIG. 10, cache SSD 505-1 is shown as including ports 805 and 810. Cache SSD 505-1 may also include cache policy 1005, SPU operations 1010 and DPU operations 1015, and cache namespaces 1020-1, 1020-2, and 1020-3. Cache namespaces 1020-1, 1020-2, and 1020-3 represent different ways to organize data moving through cache SSD 505-1. For example, data relating to one application may go through one cache namespace, data for another application may go through a second cache namespace, and so on. A single cache namespace may support multiple applications, and a single application might use multiple cache namespaces to manage its data: embodiments of the inventive concept are intended to cover all such variations.

(69) Cache policy 1005 represents rules by which data may be stored in and evicted from cache SSD 505-1. There may be multiple cache policies, and different applications may specify different cache policies. For example, one application may prefer to use a Least Recently Used (LRU) cache policy, whereas another application may prefer to use a Least Frequently Used (LFU) cache policy. Other cache policies may also be used.

(70) SPU operations 1010 and DPU operations 1015 describe operations that may be performed by SPUs 510-1 and 510-2 and DPUs 515-1 and 515-2 of FIG. 5. Example operations that may be performed by SPU 510-1 and 510-2 of FIG. 5 may include a NOP (no operation), Insert (to insert

partial data into a processing unit buffer), Write Once, XOR (exclusive OR), Extended Op (for example, data reconstruction from erasure coding information), and compression. But these operations are merely exemplary, and other operations may also be performed. More obvious examples include reading partial data from a processing unit buffer, dividing data into portions for erasure coding (such as spreading data across multiple storage devices, as well as generating parity information for data reconstruction in case a storage device fails), decompression, and so on.

(71) FIGS. 11A-11C shows details of cache namespaces **1020-1**, **1020-2**, and **1020-3** of FIG. 10. In FIG. 11A, cache namespace **1020-1** is shown. Cache namespace **1020-1** may include submission queues **1105-1** and **1110-1**, completion queues **1105-2**, and **1110-2**, SPU **510-1**, DPU **515-1** (which may be omitted), persistent memory **815**, cache manager **1115**, processing unit (PU) buffer **1120**, and multiplexer **1125**. SPU **510-1**, DPU **515-1**, and persistent memory **815** have been described above with references to FIGS. 5 and 8, and will not be repeated here.

(72) Submission queues **1105-1** and **1110-1** and completion queues **1110-1** and **1110-2** may be used to send and receive data from various sources. Thus, submission queue **1105-1** may include requests submitted to cache namespace **1020-1**: for example, read request **705** or write request **710** of FIG. 7 that may originate from processor **205** of FIG. 5. Completion queue **1110-1** may be used to send response **715** back to processor **205** of FIG. 5. Similarly, submission queue **1105-2** may include requests submitted to KV-SSDs **520-1** and/or **520-2** of FIG. 5 by cache namespace **1020-1**: for example, read request **720** or write request **725** of FIG. 7, and completion queue **1110-2** may be used to receive response **730** back from KV-SSDs **520-1** and/or **520-2** of FIG. 5.

(73) Submission queue **1105-1** and completion queue **1110-1** may form one queue pair, and submission queue **1105-2** and completion queue **1110-2** may form another queue pair. Each queue pair may be associated with a port on cache SSD **505-1** of FIG. 5. Thus, for example, submission queue **1105-1** and completion queue **1110-1** may be associated with port **805** of FIG. 8, and submission queue **1105-2** and completion queue **1110-2** may be associated with port **810** of FIG. 8.

(74) Queue pairs such as those shown in FIG. 11A are not necessarily physically present; they may be merely symbolic. That is, there may not be a separate physical queue pair for each cache namespace within cache SSD **505-1** of FIG. 5, but each queue pair may be more a “virtual” queue pair for the associated cache namespace. Which “virtual” queue pair a particular message belongs to may be identified by other data in the request: for example, the application issuing the request (for requests received via submission queue **1105-1**) or data in the result (for results received via completion queue **1110-2**) that identifies the request to which the result is responsive.

(75) Cache manager **1115** may use cache policy **1005** of FIG. 10 along with information from SPU **510-1** and/or DPU **515-1** (if included in cache namespace **1020-1**) to determine how to handle a particular request received from processor **205** of FIG. 5 via submission queue **1105-1**. Thus, cache manager **1115** may access both submission queue **1105-1** and completion queue **1110-1** (to access requests **705** and/or **710** of FIG. 7 received from processor **205** of FIG. 5 and to send result **715** of FIG. 7 back to processor **205** of FIG. 5. Cache manager **1115** may also communicate with both SPU **510-1** and DPU **515-1** to determine how to handle the request. Further, since handling a particular request may involve sending further requests to KV-SSDs **520-1** and/or **520-2** of FIG. 5, cache manager **1115** may also access submission queue **1105-2** (to send requests **720** and/or **725** of FIG. 7 to KV-SSDs **520-1** and/or **520-2** of FIG. 5) and completion queue **1110-2** (to access result **730** of FIG. 7 received from KV-SSDs **520-1** and/or **520-2** of FIG. 5).

(76) PU buffer **1120** may be used to disassemble data (when writing data to multiple KV-SSDs **520-1** and **520-2** of FIG. 5: for example, when data is being split across multiple storage devices) or to reassemble data (when reading data from multiple KV-SSDs **520-1** and **520-2** of FIG. 5: for example, when data was previously split across multiple storage devices. Finally, multiplexer **1125** may be used to direct data to various destinations, based on its intended use.

(77) FIG. 11B shows a flow diagram for handling write request **710** of FIG. 7 from processor **205** of FIG. 5. In FIG. 11B, the various components within cache namespace **1020-1** have been

rearranged to show data flow more simply, but this rearrangement is simply for simpler understanding of FIG. 11B. In FIG. 11B, both data and control are used by cache namespace **1020-1**. Control flow is depicted using solid lines, and data flow is depicted using dashed lines.

(78) When write request **710** of FIG. 7 is accessed from submission queue **1105-1**, write request **710** of FIG. 7 may be delivered to cache manager **1115**, and the data is sent to PU buffer **1120**, and thence to multiplexer **1125**. Cache manager **1115** may then determine, using cache policy **1005** of FIG. 10, where the data should be sent from multiplexer **1125**. If the data is to be stored within cache namespace **1020-1** (and not written to KV-SSDs **520-1** and/or **520-2** of FIG. 5, then the data may be directed to persistent memory **815** (the data may also be sent to flash storage **820** of FIG. 8 in addition to or instead of persistent memory **815**). If the data is to be stored on KV-SSDs **520-1** and/or **520-2** of FIG. 5, then the data may be directed to DPU **515-1**. (Note that these alternatives are not “either-or” alternatives: the data may be sent to both persistent memory **815** and ultimately to KV-SSDs **520-1** and/or **520-2** of FIG. 5.)

(79) DPU **515-1** may perform any appropriate pre-storage operations handled by DPU **515-1**. For example, if some acceleration function is to be applied to the data before storage, DPU **515-1** may execute this acceleration function. The data may then be delivered to SPU **510-1**, which may perform storage-related operations. For example, SPU **510-1** may encode, encrypt, and/or compress the data; SPU **510-1** may add Error Correcting Codes to the data; and SPU **510-1** may divide the data into smaller portions to perform erasure coding (again, these are not alternatives: SPU **510-1** may perform any desired combination of storage-related operations on the data prior to storage.) SPU **510-1** (and/or cache manager **1115**) may then generate the appropriate commands to store the data on KV-SSDs **520-1** and/or **520-2** of FIG. 5 (reflected by the convergence of the lines representing control and data flow to submission queue **1110-1**. Such command generation may include, for example, generating a KV command to store an object with a particular key on KV-SSDs **520-1** and/or **520-2** of FIG. 5. This key may be generated in any manner desired: for example, as a hash of the name of the file representing the data or some other identifier associated with the data (such as the LBA assigned to the data by processor **205** of FIG. 5). By placing the appropriate command in submission queue **1110-1**, the data may then be stored on KV-SSDs **520-1** and/or **520-2** of FIG. 5. (KV-SSDs **520-1** and **520-2** of FIG. 5 may handle the KV commands using conventional KV-SSD storage techniques.) The keys used to store the objects containing the data may be stored for later use: persistent memory **815** or mapping **920** of FIG. 9 are two possible locations to store such key information.

(80) Note that the above discussion said that the data may be divided into portions, so that different portions may be stored on different KV-SSDs **520-1** and/or **520-2** of FIG. 5. (Of course, even with the data being divided into different portions, at least some of the different portions might end up being stored on the same KV-SSD, much like how a file that is too large to fit into a RAID stripe across the available disks may store data in multiple stripes that span all of the available disks.) While KV storage devices may expect each object stored on a single KV storage device to have a unique key, it is not necessarily true that keys are unique across multiple KV storage devices. That is, the same key might be used to store data on each of KV-SSDs **520-1** and **520-2** of FIG. 5. SPU **510-1** (and cache manager **1115**) may leverage this fact to keep the keys that identify objects including data that is part of the same file related. In other words, the same key may be used to store one portion of data on KV-SSD **520-1** of FIG. 5, a second portion of data on KV-SSD **520-2** of FIG. 5, and so on (up to the number of KV storage devices in storage server **120** of FIG. 5).

(81) Alternatively, SPU **510-1** (and cache manager **1115**) may assign different (but still potentially related) keys to portions of data stored on different KV storage devices. For example, after hashing the file name or other data (such as the LBA), SPU **510-1** (and cache manager **1115**) might append an index to each key to identify the order in which the portions of data should be reassembled upon read request **705** of FIG. 7. (Cache namespace **1020-1** may track such ordering information regardless of what methodology is used to generate the keys, but there is no harm in redundant

approaches.) Thus, if the key assigned to the data was the string “object key”, the key used for the first portion of the data might be “object key_1”, the key used for the second portion of the data might be “object key_2”, and so on.

(82) There are a number of reasons why SPU **510-1** (and cache manager **1115**) might divide the data into portions. Aside from erasure coding, it might be that the total data is too large to fit in a single object on KV-SSDs **520-1** and/or **520-2** of FIG. 5 (assuming that KV-SSDs **520-1** and/or **520-2** of FIG. 5 have limits on the size of an individual object). Or the size of the data may be too large to fit into a single KV-SSD (for example, the largest KV-SSD may be 1 terabyte in size, but the data requires 2 terabytes to store). But assuming the data would fit on a single KV-SSD, that the KV-SSD is capable of internally managing object size (and divide large data into smaller chunks of data that would fit into objects whose size may be limited), and that there is no other reason to divide the data into portions, SPU **510-1** (and cache manager **1115**) may just send the entire object to KV-SSDs **520-1** and/or **520-2** of FIG. 5 without first dividing the data into portions.

(83) It is also possible that SPU **510-1** (and cache manager **1115**) might store the same data in multiple places (perhaps for redundancy). The same key may be used to store the same data on more than one KV storage device. When such embodiments of the inventive concept are used, when processing read request **705** of FIG. 7 for the data, cache namespace **1020-1** might send a request to read the data (using the appropriate key) to any number of KV-SSDs **520-1** and/or **520-2** of FIG. 5 (that is, cache namespace **1020-1** might request the data from only one KV storage device, from two KV storage devices, and so on up to all available KV storage devices). KV storage devices that store the data may then return the requested data; KV storage devices that do not store the data may return messages indicating the data was not found. Cache namespace **1020-1** might then return the data from any of the KV storage devices (since any KV storage device that stores the data should store the same data), and may discard any other responses (both successful data retrievals and negative responses).

(84) Returning to the control flow, eventually KV-SSDs **520-1** and/or **520-2** of FIG. 5 should place response **730** of FIG. 7 in completion queue **1110-2**. This result may be retrieved from completion queue **1110-2** by cache manager **1115** and result **715** of FIG. 7 may be placed in completion queue **1105-2** for later use by processor **205** of FIG. 5. Result **715** of FIG. 7 may just be an unedited forwarding of result **730** of FIG. 7, or it may be an edited forwarding of result **730** of FIG. 7, or it may be an entirely new message, albeit perhaps returning the same general information—namely, that the data was successfully stored.

(85) Although the above discussion suggests that all storage-related processing of the data is performed by SPU **510-1** within SPU **510-1**, some embodiments of the inventive concept may perform such processing elsewhere (although perhaps under instruction from SPU **510-1** and/or cache manager **1115**). For example, PU buffer **1120** may be a buffer which may store the data, and SPU **510-1** (and cache manager **1115**) may process the data while in PU buffer **1120**. Alternatively, PU buffer **1120**, while represented as separate from SPU **510-1** in FIGS. **11A-11C**, may actually be implemented within SPU **510-1**, thus enabling SPU **510-1** to perform storage-related processing of the data as described.

(86) Although not explicitly described above, at any point (once the data is within cache namespace **1020-1** and protected against loss due to power failure) cache namespace **1020-1** may place result **715** of FIG. 7 in completion queue **1105-2**, providing an early acknowledgment of write request **710** of FIG. 7 to processor **205** of FIG. 5. Alternatively, cache namespace **1020-1** might not place result **715** of FIG. 7 in completion queue **1105-2** until after receiving result **730** of FIG. 7 from KV-SSDs **520-1** and/or **520-2** of FIG. 5 via completion queue **1110-2**. Embodiments of the inventive concept are intended to cover both such variations.

(87) Note that in the above discussion, how KV-SSDs **520-1** and/or **520-2** of FIG. 5 are selected to store the data is not discussed. Cache namespace **1020-1** may use any desired methodology to select which KV-SSDs store the data. Possible methodologies that may be used include: Random

selection from among available KV-SSDs. Round robin selection of KV-SSDs (that is, identifying the KV-SSDs in a particular sequence, and selecting the next KV-SSD from the sequence, restarting from the first KV-SSD after the last KV-SSD in the sequence has been used). Balancing the amount of data stored on each KV-SSD (in other words, selecting the KV-SSD with the smallest amount of stored data). Balancing the percentage of data stored on each KV-SSD (in other words, selecting the KV-SSD that has the smallest amount of stored data relative to its total capacity). Selecting a KV-SSD to satisfy Quality of Service (QoS) requirements of the application storing the data (QoS requirements may include minimum or maximum latency of the storage device, minimum or maximum bandwidth of the storage device, encryption offered by the storage device, and isolation from data of another application: U.S. patent application Ser. No. 15/227,955, filed Aug. 3, 2016, now U.S. Pat. No. 10,437,486, issued Oct. 8, 2019, which claims the benefit of U.S. Provisional Patent Application Ser. No. 62/341,063, filed May 24, 2016, both of which are incorporated by reference herein for all purposes, describes other possible QoS requirements that may be applied).

(88) In addition, multiple storage devices may be selected. As discussed above, cache namespace **1020-1** may divide the data into portions and store different portions on different storage devices. Cache namespace **1020-1** may also select multiple storage devices to store the data for redundancy purposes (redundancy may also be a QoS requirement of the application).

(89) Once storage devices have been selected to store the data (or portions thereof), cache namespace **1020-1** may keep track of which storage devices have been selected. This information may be stored in persistent memory **815**, mapping **920** of FIG. **9**, or any other desired location. This information (which storage devices store the data) may also be stored with the key(s) used to store the data on the storage devices. But note that neither which storage devices, nor which keys are used to store the data, are absolutely necessary. As discussed above, the keys may be generated from other information (such as the file name or the LBA of the data), permitting their recovery. And while storing which storage devices store the data is helpful, it is also possible to simply attempt to retrieve the data from all storage devices. The storage devices that store the data (or portions thereof) may respond with what they store, while storage devices that do not store the data may respond negatively (since the key provided by cache namespace **1020-1** should not match up with a key storing data on those storage devices). Storing information about which storage devices store the data may simply avoid querying devices that do not store the data.

(90) Note too that if the data is divided into portions, that information may be stored with the information about which storage devices store the data and/or the keys used as well. Storing information about how the data is divided may expedite reassembly of the complete data. But if the portions of the data include information about their order (for example, by including a header that identifies that a particular portion of data is portion *i* of *n*), then cache namespace **1020-1** may not need to store this information either.

(91) FIG. **11C** shows a flow diagram for handling read request **705** of FIG. **7** from processor **205** of FIG. **5**. In FIG. **11C**, the various components within cache namespace **1020-1** have been rearranged to show data flow more simply, but this rearrangement is simply for simpler understanding of FIG. **11C**. In FIG. **11C**, both data and control is used by cache namespace **1020-1**. Control flow is depicted using solid lines, and data flow is depicted using dashed lines.

(92) When read request **705** of FIG. **7** is accessed from submission queue **1105-1**, read request **705** of FIG. **7** may be delivered to cache manager **1115**. Cache manager **1115** may then generate appropriate KV commands (there may be one or more KV commands, depending on how the data was stored) to retrieve the object storing the requested data, and place these KV commands in submission queue **1110-1**. The appropriate KV commands may be generated in FIG. **11C** similarly to the KV command discussed above with reference to FIG. **11B**, except that the KV command may be to retrieve data rather than to store data. The appropriate key(s) may be generated (for example, by hashing the file name requested or the LBA of the data), with indices added if appropriate to how the objects were stored.

(93) When KV-SSDs **520-1** and/or **520-2** of FIG. 5 return the requested data, results **730** of FIG. 7 may be placed in completion queue **1110-2**. SPU **510-1**, in coordination with cache manager **1115**, may then be notified and the data loaded into PU buffer **1120**. SPU **510-1** may then perform any storage-related processing on the data in PU buffer **1120**. For example, if the data was stored in multiple objects, SPU **510-1** may reassemble the original data by placing the portions in the correct order in PU buffer **1120**. If the data was compressed, encoded, and/or encrypted, SPU **510-1** may perform the appropriate decompression/decoding/decryption. If Error Correcting Codes were added to the data, SPU **510-1** may use the Error Correcting Codes to verify that the data was correctly returned and, if not, attempt to correct the errors (or notify processor **205** of FIG. 5 that the data was not successfully retrieved). As with SPU **510-1** in FIG. 11B, SPU **510-1** may perform any desired combination of operations appropriate to the data.

(94) Once SPU **510-1** has completed any storage-related operations on the data, the data may be delivered to DPU **515-1**, which may then perform any additional functions on the data appropriate to DPU **515-1** (in coordination with cache manager **1115**). For example as discussed above, the data might involve a large data set with many entries, but processor **205** of FIG. 5 may only be interested in knowing how many entries in the data set meet a particular condition. DPU **515-1** may perform the appropriate query on the data and instead return the result of the query (rather than the entire data set).

(95) After DPU **515-1** of FIG. 1 has finished processing the data, the data may be delivered to multiplexer **1125**. Multiplexer **1125**, in coordination with cache manager **1115**, may then direct the data to either persistent memory **815** or to completion queue **1105-2** for return to processor **205** of FIG. 5. (As with FIG. 11B, multiplexer **1125** may direct the data to both persistent memory **815** and to completion queue **1105-2**; the choice is not an “either-or” choice.) Once result **715** of FIG. 7 is placed in completion queue **1105-2**, processor **205** of FIG. 5 may access result **715** of FIG. 7 and process it accordingly.

(96) FIGS. 12A-12B show a flowchart of an example procedure for cache SSDs **505-1** and **505-2** of FIG. 5 to process write request **710** of FIG. 7 from compute server **105** of FIG. 1, according to an embodiment of the inventive concept. In FIG. 12A, at block **1205**, cache SSD **505-1** and/or **505-2** of FIG. 5 may receive write request **710** of FIG. 7 from processor **205** of FIG. 5. Cache SSDs **505-1** and **505-2** may receive write request **710** of FIG. 7 via submission queue **1105-1** of FIG. 11B over port **805** of FIG. 8. At block **1210**, DPU **515-1** may apply an acceleration function to the data. Block **1210** is optional, as shown by dashed line **1215**. At block **1220**, the data may be stored in cache SSD **505-1** and/or **505-2** of FIG. 5. Such storage may be within flash storage **820** of FIG. 8, persistent memory **815** of FIG. 11B, PU buffer **1120** of FIG. 11B, or any other desired location (as long as the data is protected against lost due to power failure). At block **1225**, cache namespace **1020-1** of FIG. 11B may send an early acknowledgment that the data was written by placing result **715** of FIG. 7 in completion queue **1105-1** of FIG. 11B, which may be delivered to processor **205** of FIG. 5 via completion queue **1105-2** of FIG. 11B over port **805** of FIG. 8. Blocks **1220** and **1225** are variously optional, as shown by dashed lines **1230** and **1235**.

(97) At block **1240** (FIG. 12B), SPU **510-1** of FIG. 11B may determine appropriate KV-SSDs **520-1** and/or **520-2** of FIG. 5 to store the data. KV-SSDs **520-1** and/or **520-2** may be intended to store the entirety of the data, or they may be intended to store only one portion (of some number of portions) of the data. As discussed above with reference to FIG. 11B, any desired methodology may be used to select the KV-SSDs to store the data (or portions thereof). At block **1245**, SPU **510-1** of FIG. 11B may disassemble the data into portions. (SPU **510-1** of FIG. 11 may also perform other storage-related operations at this time, as described above with reference to FIG. 11B.) Of course, if the data is being stored in a single object within a single KV-SSD, then block **1245** may be omitted, as shown by dashed line **1250**. At block **1255**, SPU **510-1** may generate the appropriate KV commands to send to KV-SSDs **520-1** and/or **520-2** of FIG. 5, and at block **1260** these KV commands may be sent to KV-SSDs **520-1** and/or **520-2** of FIG. 5 by being placed in submission

queue **1110-1** of FIG. **11B**, which may be delivered to KV-SSDs **520-1** and/or **520-2** of FIG. **5** via port **810** of FIG. **8**. Cache namespace **1020-1** of FIG. **11B** may also store information about how the data is stored on KV-SSDs **520-1** and/or **520-2** of FIG. **5**: for example, by identifying which KV-SSDs store the data or portions of the data, what keys are used to store the data, and which portions of the data each KV-SSD stores. But as discussed above with reference to FIG. **11B**, depending on how the data is stored on KV-SSDs **520-1** and/or **520-2** of FIG. **5**, none of this information may be necessary to be able to retrieve and reassemble the data from KV-SSDs **520-1** and/or **520-2** of FIG. **5**.

(98) At block **1265**, cache manager **1115** of FIG. **11B** may access results **730** of FIG. **7** from completion queue **1110-2** of FIG. **11B**, which may be received from KV-SSDs **520-1** and/or **520-2** of FIG. **5** via port **810** of FIG. **8**. These results **730** of FIG. **7** may originate from the KV-SSDs **520-1** and/or **520-2** that were instructed to store some or all of the data. At block **1270**, cache manager **1115** of FIG. **11B** may place result **715** of FIG. **7** in completion queue **1105-2** of FIG. **11B**, to inform processor **205** of FIG. **5** over port **805** of FIG. **8** that the data has been successfully stored. Block **1270** is optional, as shown by dashed line **1275**; but at least one of blocks **1225** (FIG. **12A**) and **1270** should be performed, or else processor **205** of FIG. **5** would not be notified that the data was successfully stored.

(99) FIGS. **13A-13B** show a flowchart of an example procedure for cache SSDs **505-1** and **505-2** of FIG. **5** to process read request **705** of FIG. **7** from compute server **105** of FIG. **1**, according to an embodiment of the inventive concept. In FIG. **13A**, at block **1305**, cache SSDs **505-1** and/or **505-2** of FIG. **5** may receive read request **705** of FIG. **7** from processor **205** of FIG. **5**. Cache SSDs **505-1** and **505-2** may receive read request **705** of FIG. **7** via submission queue **1105-1** of FIG. **11C** over port **805** of FIG. **8**. At block **1310**, cache SSDs **505-1** and/or **505-2** of FIG. **5** may check to see if the data requested in read request **705** of FIG. **7** is stored locally (perhaps in persistent memory **815** of FIG. **11C** or flash storage **825** of FIG. **8**). If the data requested in read request **705** of FIG. **7** is stored locally, then at block **1315** cache SSDs **505-1** and/or **505-2** of FIG. **5** may read the data from local storage. At block **1320**, DPUs **515-1** and/or **515-2** of FIG. **5** may apply an acceleration function on the data. Block **1320** is optional, as shown by dashed line **1325**. Finally, at block **1330**, cache SSDs **505-1** and/or **505-2** of FIG. **5** may return the data to processor **205** of FIG. **5** by placing result **715** of FIG. **7** in completion queue **1105-2** of FIG. **11C**, which may be delivered to processor **205** of FIG. **5** via port **805** of FIG. **8**.

(100) On the other hand, if the data requested in read request **705** of FIG. **7** is not stored on cache SSDs **505-1** and/or **505-2**, then at block **1335** (FIG. **13B**), SPU **510-1** of FIG. **11C** (possibly working in conjunction with cache manager **1115** of FIG. **11C**) may generate KV requests to read the data from KV-SSDs **520-1** and/or **520-2** of FIG. **5**. At block **1340**, KV-SSDs **520-1** and/or **520-2** of FIG. **5** that store the data may be identified. Block **1340** may be omitted, as shown by dashed line **1345**, although in that case cache namespace **1020-1** of FIG. **11C** may need to send requests to every KV-SSD in storage server **120** of FIG. **5** to retrieve the data.

(101) At block **1350**, SPU **510-1** of FIG. **11C** (and cache manager **1115** of FIG. **11C**) may send read request **720** of FIG. **7** to KV-SSDs **520-1** and/or **520-2** of FIG. **5** by placing read requests **720** of FIG. **7** in submission queue **1110-1** of FIG. **11C**, which may be delivered to KV-SSDs **520-1** and/or **520-2** of FIG. **5** over port **810** of FIG. **8**. At block **1355**, cache namespace **1020-1** of FIG. **11C** may receive results **730** of FIG. **7** from KV-SSDs **520-1** and/or **520-2** of FIG. **5**: results **730** of FIG. **7** may be received via completion queue **1120-2** of FIG. **11C** over port **810** of FIG. **8**. At block **1360**, SPU **510-1** of FIG. **11C** (and cache manager **1115** of FIG. **11C**) may reassemble the data from portions received from KV-SSDs **520-1** and/or **520-2** of FIG. **5**. Obviously, if the data is received in its entirety from a KV-SSD, then there is no need to reassemble the data: block **1360** may therefore be omitted as shown by dashed line **1365**. Processing may then return to block **1320** of FIG. **13A** to complete delivery of the requested data to processor **205** of FIG. **5**.

(102) Note that it may occur that cache SSDs **505-1** and/or **505-2** of FIG. **5** have some of the data,

but not necessarily all of the data, requested in read request **705** of FIG. **7**. For example, the data might span more than one “cache line” within cache SSDs **505-1** and/or **505-2** of FIG. **5**: if one “cache line” is evicted but another is not, the data may be partially stored in cache SSDs **505-1** and **505-2** of FIG. **5**. (Of course, as cache SSDs **505-1** and **505-2** of FIG. **5** may be SSDs themselves, the data stored thereon may not be organized into cache lines as that term is conventionally used, but the approach is similar.) In such situations, cache namespace **1020-1** of FIG. **11C** may be able to determine which data is stored on cache SSDs **505-1** and **505-2** of FIG. **5** and what data is not stored on cache SSDs **505-1** and **505-2** of FIG. **5** and send read requests **720** of FIG. **7** to retrieve only the missing data. For example, cache namespace **1020-1** may determine what portions of the data are stored on cache SSDs **505-1** and **505-2** of FIG. **5**, identify the portions of the data that are missing, identify which KV-SSDs store the missing portions, request those missing portions, and reassemble the data from both the data stored on cache SSDs **505-1** and **505-2** of FIG. **5** and the data received from KV-SSDs **520-1** and/or **520-2** of FIG. **5**.

(103) In FIGS. **12A-13B**, some embodiments of the inventive concept are shown. But a person skilled in the art will recognize that other embodiments of the inventive concept are also possible, by changing the order of the blocks, by omitting blocks, or by including links not shown in the drawings. All such variations of the flowcharts are considered to be embodiments of the inventive concept, whether expressly described or not.

(104) Embodiments of the inventive concept offer technical advantages over the prior art. In conventional systems, between executing applications, managing the operating system and file system, and controlling devices such as storage devices, the burden on the local processor may be high. Embodiments of the inventive concept may shift the burden off the local processor without introducing new bottlenecks in other parts of the system (such as another processor or memory). In addition, embodiments of the inventive concept may permit data to be stored using a storage methodology without the local processor being aware of the actual storage methodology. Thus, for example, the local processor may send data read and write requests using a block-based approach even though the actual storage devices may use a Key-Value storage system.

(105) The following discussion is intended to provide a brief, general description of a suitable machine or machines in which certain aspects of the inventive concept may be implemented. The machine or machines may be controlled, at least in part, by input from conventional input devices, such as keyboards, mice, etc., as well as by directives received from another machine, interaction with a virtual reality (VR) environment, biometric feedback, or other input signal. As used herein, the term “machine” is intended to broadly encompass a single machine, a virtual machine, or a system of communicatively coupled machines, virtual machines, or devices operating together. Exemplary machines include computing devices such as personal computers, workstations, servers, portable computers, handheld devices, telephones, tablets, etc., as well as transportation devices, such as private or public transportation, e.g., automobiles, trains, cabs, etc.

(106) The machine or machines may include embedded controllers, such as programmable or non-programmable logic devices or arrays, Application Specific Integrated Circuits (ASICs), embedded computers, smart cards, and the like. The machine or machines may utilize one or more connections to one or more remote machines, such as through a network interface, modem, or other communicative coupling. Machines may be interconnected by way of a physical and/or logical network, such as an intranet, the Internet, local area networks, wide area networks, etc. One skilled in the art will appreciate that network communication may utilize various wired and/or wireless short range or long range carriers and protocols, including radio frequency (RF), satellite, microwave, Institute of Electrical and Electronics Engineers (IEEE) 802.11, Bluetooth®, optical, infrared, cable, laser, etc.

(107) Embodiments of the present inventive concept may be described by reference to or in conjunction with associated data including functions, procedures, data structures, application programs, etc. which when accessed by a machine results in the machine performing tasks or

defining abstract data types or low-level hardware contexts. Associated data may be stored in, for example, the volatile and/or non-volatile memory, e.g., RAM, ROM, etc., or in other storage devices and their associated storage media, including hard-drives, floppy-disks, optical storage, tapes, flash memory, memory sticks, digital video disks, biological storage, etc. Associated data may be delivered over transmission environments, including the physical and/or logical network, in the form of packets, serial data, parallel data, propagated signals, etc., and may be used in a compressed or encrypted format. Associated data may be used in a distributed environment, and stored locally and/or remotely for machine access.

(108) Embodiments of the inventive concept may include a tangible, non-transitory machine-readable medium comprising instructions executable by one or more processors, the instructions comprising instructions to perform the elements of the inventive concepts as described herein.

(109) The various operations of methods described above may be performed by any suitable means capable of performing the operations, such as various hardware and/or software component(s), circuits, and/or module(s). The software may comprise an ordered listing of executable instructions for implementing logical functions, and may be embodied in any “processor-readable medium” for use by or in connection with an instruction execution system, apparatus, or device, such as a single or multiple-core processor or processor-containing system.

(110) The blocks or steps of a method or algorithm and functions described in connection with the embodiments disclosed herein may be embodied directly in hardware, in a software module executed by a processor, or in a combination of the two. If implemented in software, the functions may be stored on or transmitted over as one or more instructions or code on a tangible, non-transitory computer-readable medium. A software module may reside in Random Access Memory (RAM), flash memory, Read Only Memory (ROM), Electrically Programmable ROM (EPROM), Electrically Erasable Programmable ROM (EEPROM), registers, hard disk, a removable disk, a CD ROM, or any other form of storage medium known in the art.

(111) Having described and illustrated the principles of the inventive concept with reference to illustrated embodiments, it will be recognized that the illustrated embodiments may be modified in arrangement and detail without departing from such principles, and may be combined in any desired manner. And, although the foregoing discussion has focused on particular embodiments, other configurations are contemplated. In particular, even though expressions such as “according to an embodiment of the inventive concept” or the like are used herein, these phrases are meant to generally reference embodiment possibilities, and are not intended to limit the inventive concept to particular embodiment configurations. As used herein, these terms may reference the same or different embodiments that are combinable into other embodiments.

(112) The foregoing illustrative embodiments are not to be construed as limiting the inventive concept thereof. Although a few embodiments have been described, those skilled in the art will readily appreciate that many modifications are possible to those embodiments without materially departing from the novel teachings and advantages of the present disclosure. Accordingly, all such modifications are intended to be included within the scope of this inventive concept as defined in the claims.

(113) Embodiments of the inventive concept may extend to the following statements, without limitation:

(114) Statement 1. An embodiment of the inventive concept includes a Solid State Drive (SSD), comprising: a first port to receive first read requests and first write requests from a host; a second port to send second read requests and second write requests to a second storage device; flash storage for data; and an SSD controller to process the first read requests to generate the second read requests and to process the first write requests to generate the second write requests, wherein the SSD is used as a cache for data stored on the second storage device.

(115) Statement 2. An embodiment of the inventive concept includes the SSD according to statement 1, wherein: the first port includes a Peripheral Component Interconnect Express (PCIe)

port; and the second port includes an Ethernet port.

(116) Statement 3. An embodiment of the inventive concept includes the SSD according to statement 1, wherein the second storage device includes a Key-Value SSD (KV-SSD).

(117) Statement 4. An embodiment of the inventive concept includes the SSD according to statement 3, wherein the KV-SSD is remote from the SSD.

(118) Statement 5. An embodiment of the inventive concept includes the SSD according to statement 3, wherein: the first read requests and the first write requests from the host include PCIe read requests and PCIe write requests from the host; and the SSD controller is operative to send key-value read requests and key-value write requests to the KV-SSD.

(119) Statement 6. An embodiment of the inventive concept includes the SSD according to statement 1, further comprising a storage-related processing unit (SPU).

(120) Statement 7. An embodiment of the inventive concept includes the SSD according to statement 6, wherein the SPU is operative to manage storage of data on the second storage device.

(121) Statement 8. An embodiment of the inventive concept includes the SSD according to statement 6, wherein the SPU includes at least one of a general purpose processor, a Field Programmable Gate Array (FPGA), an Application-Specific Integrated Circuit (ASIC), a Graphic Processing Unit (GPU), and a General Purpose Graphics Processing Unit (GPGPU).

(122) Statement 9. An embodiment of the inventive concept includes the SSD according to statement 1, further comprising a general data processing unit (DPU).

(123) Statement 10. An embodiment of the inventive concept includes the SSD according to statement 9, wherein the DPU is operative to execute instructions of an application.

(124) Statement 11. An embodiment of the inventive concept includes the SSD according to statement 10, wherein the application is off-loaded from the host.

(125) Statement 12. An embodiment of the inventive concept includes the SSD according to statement 1, wherein the SSD is organized into at least one cache namespace.

(126) Statement 13. An embodiment of the inventive concept includes the SSD according to statement 12, wherein the SSD includes at least one first queue and at least one second queue for the at least one cache namespace.

(127) Statement 14. An embodiment of the inventive concept includes the SSD according to statement 13, wherein the SSD includes the at least one first queue and the at least one second queue for each namespace in the at least one cache namespace.

(128) Statement 15. An embodiment of the inventive concept includes the SSD according to statement 13, wherein: the at least one first queue includes at least one first submission queue and at least one first completion queue; and the at least one second queue includes at least one second submission queue and at least one second completion queue.

(129) Statement 16. An embodiment of the inventive concept includes the SSD according to statement 13, wherein: the at least one first queue is associated with the first port; and the at least one second queue is associated with the second port.

(130) Statement 17. An embodiment of the inventive concept includes the SSD according to statement 16, wherein the SSD controller is operative to receive the first read requests and the first write requests in the at least one first queue associated with the first port and to send the second read requests and the second write requests using the at least one second queue associated with the second port.

(131) Statement 18. An embodiment of the inventive concept includes the SSD according to statement 13, wherein the at least one cache namespace further includes: a cache manager to manage requests received via the at least one first queue and to issue requests via the at least one second queue; and a persistent memory to store data based on the requests.

(132) Statement 19. An embodiment of the inventive concept includes the SSD according to statement 18, wherein the at least one cache namespace further includes a processing unit (PU) buffer for data assembly and disassembly.

(133) Statement 20. An embodiment of the inventive concept includes the SSD according to statement 1, wherein the SSD is operative to respond to the first read requests based at least in part on the data in the flash storage.

(134) Statement 21. An embodiment of the inventive concept includes the SSD according to statement 1, wherein the SSD is operative to respond to the first write requests based at least in part on the data in the flash storage.

(135) Statement 22. An embodiment of the inventive concept includes the SSD according to statement 1, wherein the SSD controller is operative to send the second read requests and the second write requests to a plurality of second storage devices.

(136) Statement 23. An embodiment of the inventive concept includes the SSD according to statement 22, wherein the SSD controller manages erasure coding of the second read requests and the second write requests on the plurality of second storage device.

(137) Statement 24. An embodiment of the inventive concept includes a system, comprising: a processor executing an operating system; a memory storing data for the operating system; a cache SSD, including: a first port to receive first read requests and first write requests from the processor; a second port to send second read requests and second write requests to a second storage device; flash storage for data; and an SSD controller to process the first read requests to generate the second read requests and to process the first write requests to generate the second write requests, and the second storage device, wherein the SSD is used as a cache for data stored on the second storage device.

(138) Statement 25. An embodiment of the inventive concept includes the system according to statement 24, wherein the second storage device includes a Key-Value SSD (KV-SSD).

(139) Statement 26. An embodiment of the inventive concept includes the system according to statement 25, wherein the KV-SSD is remote from the SSD.

(140) Statement 27. An embodiment of the inventive concept includes the system according to statement 25, wherein: the first read requests and the first write requests from the host include PCIe read requests and PCIe write requests from the host; and the SSD controller is operative to send key-value read requests and key-value write requests to the KV-SSD.

(141) Statement 28. An embodiment of the inventive concept includes the system according to statement 25, wherein the KV-SSD includes an SPU to manage storage of data on the KV-SSD.

(142) Statement 29. An embodiment of the inventive concept includes the system according to statement 24, wherein the cache SSD further includes a storage-related processing unit (SPU).

(143) Statement 30. An embodiment of the inventive concept includes the system according to statement 29, wherein the SPU is operative to manage storage of data on the second storage device.

(144) Statement 31. An embodiment of the inventive concept includes the system according to statement 30, wherein the SPU is operative to perform an acceleration function on the data on the second storage device.

(145) Statement 32. An embodiment of the inventive concept includes the system according to statement 31, wherein the acceleration function is at least one of encryption, compression, decryption, or decompression.

(146) Statement 33. An embodiment of the inventive concept includes the system according to statement 29, wherein the SPU includes at least one of a general purpose processor, a Field Programmable Gate Array (FPGA), an Application-Specific Integrated Circuit (ASIC), a Graphic Processing Unit (GPU), and a General Purpose Graphics Processing Unit (GPGPU).

(147) Statement 34. An embodiment of the inventive concept includes the system according to statement 24, wherein the SSD controller is operative to send the second read requests and the second write requests to a plurality of second storage devices.

(148) Statement 35. An embodiment of the inventive concept includes the system according to statement 24, further comprising a general data processing unit (DPU).

(149) Statement 36. An embodiment of the inventive concept includes a method comprising:

receiving a first write request for a data from a host at a Solid State Drive (SSD) over a first port of the SSD; sending a second write request for the data from the SSD to a second storage device over a second port of the SSD; and receiving a write result from the second storage device at the SSD over the second port of the SSD; and sending the write result from the SSD to the host over the first port of the SSD.

(150) Statement 37. An embodiment of the inventive concept includes the method according to statement 36, wherein receiving a first write request for a data from a host at a Solid State Drive (SSD) over a first port of the SSD includes receiving the first write request for the data from the host at the SSD via a first queue over the first port of the SSD; and sending the write result from the SSD to the host over the first port of the SSD includes sending the write result from the SSD to the host via the first queue over the first port of the SSD.

(151) Statement 38. An embodiment of the inventive concept includes the method according to statement 37, wherein the first queue includes queue pair including a first submission queue and a first completion queue.

(152) Statement 39. An embodiment of the inventive concept includes the method according to statement 36, wherein: sending a second write request for the data from the SSD to a second storage device over a second port of the SSD includes sending the second write request for the data from the SSD to the second storage device via a second queue over the second port of the SSD; and receiving a write result from the second storage device at the SSD over the second port of the SSD includes receiving the write result from the second storage device at the SSD via the second queue over the second port of the SSD.

(153) Statement 40. An embodiment of the inventive concept includes the method according to statement 39, wherein the second queue includes queue pair including a second submission queue and a second completion queue.

(154) Statement 41. An embodiment of the inventive concept includes the method according to statement 36, wherein sending a second write request for the data from the SSD to a second storage device over a second port of the SSD includes sending the second write request for the data from the SSD to a plurality of second storage devices over the second port of the SSD.

(155) Statement 42. An embodiment of the inventive concept includes the method according to statement 41, wherein sending the second write request for the data from the SSD to a plurality of second storage devices over the second port of the SSD includes: disassembling the data into at least two portions of the data; and sending a portion of the data from the SSD to one of the plurality of second storage devices over the second port of the SSD.

(156) Statement 43. An embodiment of the inventive concept includes the method according to statement 42, wherein sending the second write request for the data from the SSD to a plurality of second storage devices over the second port of the SSD further includes identifying the one of the plurality of second storage devices using a SPU before sending the portion of the data from the SSD to one of the plurality of second storage devices over the second port of the SSD.

(157) Statement 44. An embodiment of the inventive concept includes the method according to statement 42, wherein disassembling the data into at least two portions of the data includes disassembling the data into at least two portions of the data using a storage processing unit (SPU).

(158) Statement 45. An embodiment of the inventive concept includes the method according to statement 41, wherein receiving the write result from the second storage device at the SSD over the second port of the SSD includes receiving the write result from the plurality of second storage devices at the SSD over the second port of the SSD.

(159) Statement 46. An embodiment of the inventive concept includes the method according to statement 36, further comprising applying an acceleration function to the data before sending the second write request for the data from the SSD to the second storage device over the second port of the SSD.

(160) Statement 47. An embodiment of the inventive concept includes the method according to

statement 46, wherein applying an acceleration function to the data before sending the second write request for the data from the SSD to the second storage device over the second port of the SSD includes applying the acceleration function to the data using a general Data Processing Unit (DPU) before sending the second write request for the data from the SSD to the second storage device over the second port of the SSD.

(161) Statement 48. An embodiment of the inventive concept includes the method according to statement 36, further comprising storing the data in a flash storage of the SSD.

(162) Statement 49. An embodiment of the inventive concept includes the method according to statement 48, further comprising sending the write result from the SSD to the host over the first port of the SSD before receiving the write result from the second storage device at the SSD over the second port of the SSD.

(163) Statement 50. An embodiment of the inventive concept includes the method according to statement 36, wherein: sending a second write request for the data from the SSD to a second storage device over a second port of the SSD sending a Key-Value (KV) write request for the data from the SSD to a Key-Value SSD (KV-SSD) over the second port of the SSD; and receiving a write result from the second storage device at the SSD over the second port of the SSD includes receiving the write result from the KV-SSD at the SSD over the second port of the SSD; and

(164) Statement 51. An embodiment of the inventive concept includes the method according to statement 50, further comprising generating the KV write request from the first write request.

(165) Statement 52. An embodiment of the inventive concept includes a method, comprising: receiving a first read request for a data from a host at a Solid State Drive (SSD) over a first port of the SSD; determining whether the data is stored in a flash storage of the SSD; based at least in part on the data not being stored in the flash storage of the SSD: sending a second read request for the data from the SSD to a second storage device over a second port of the SSD; and receiving the data from the second storage device at the SSD over the second port of the SSD; and sending the data from the SSD to the host over the first port of the SSD.

(166) Statement 53. An embodiment of the inventive concept includes the method according to statement 52, wherein receiving a first read request for a data from a host at a Solid State Drive (SSD) over a first port of the SSD includes receiving the first read request for the data from the host at the SSD via a first queue over the first port of the SSD; and sending the data from the SSD to the host over the first port of the SSD includes sending the data from the SSD to the host via the first queue over the first port of the SSD.

(167) Statement 54. An embodiment of the inventive concept includes the method according to statement 53, wherein the first queue includes queue pair including a first submission queue and a first completion queue.

(168) Statement 55. An embodiment of the inventive concept includes the method according to statement 52, wherein: sending a second read request for the data from the SSD to a second storage device over a second port of the SSD includes sending the second read request for the data from the SSD to the second storage device via a second queue over the second port of the SSD; and receiving the data from the second storage device at the SSD over the second port of the SSD includes receiving the data from the second storage device at the SSD via the second queue over the second port of the SSD.

(169) Statement 56. An embodiment of the inventive concept includes the method according to statement 55, wherein the second queue includes queue pair including a second submission queue and a second completion queue.

(170) Statement 57. An embodiment of the inventive concept includes the method according to statement 52, wherein sending a second read request for the data from the SSD to a second storage device over a second port of the SSD includes sending the second read request for the data from the SSD to a plurality of second storage devices over the second port of the SSD.

(171) Statement 58. An embodiment of the inventive concept includes the method according to

statement 57, wherein receiving the data from the second storage device at the SSD over the second port of the SSD includes receiving the data from the plurality of second storage devices at the SSD over the second port of the SSD.

(172) Statement 59. An embodiment of the inventive concept includes the method according to statement 58, wherein receiving the data from the plurality of second storage devices at the SSD over the second port of the SSD includes: receiving at least two portions of the data from the plurality of second storage devices at the SSD over the second port of the SSD; and reassembling the data from the at least two portions of the data.

(173) Statement 60. An embodiment of the inventive concept includes the method according to statement 59, wherein reassembling the data from the at least two portions of the data includes reassembling the data from the at least two portions of the data using a storage processing unit (SPU).

(174) Statement 61. An embodiment of the inventive concept includes the method according to statement 57, further comprising identifying the plurality of second storage devices using a SPU before sending the second read request for the data from the SSD to the plurality of second storage devices over the second port of the SSD.

(175) Statement 62. An embodiment of the inventive concept includes the method according to statement 52, further comprising applying an acceleration function to the data before sending the data from the SSD to the host over the first port of the SSD.

(176) Statement 63. An embodiment of the inventive concept includes the method according to statement 62, wherein applying an acceleration function to the data before sending the data from the SSD to the host over the first port of the SSD includes applying the acceleration function to the data using a general Data Processing Unit (DPU) before sending the data from the SSD to the host over the first port of the SSD.

(177) Statement 64. An embodiment of the inventive concept includes the method according to statement 52, further comprising, based at least in part on the data being stored in the flash storage of the SSD, reading the data from the flash storage of the SSD.

(178) Statement 65. An embodiment of the inventive concept includes the method according to statement 52, wherein: sending a second read request for the data from the SSD to a second storage device over a second port of the SSD includes sending a Key-Value (KV) read request for the data from the SSD to a Key-Value SSD (KV-SSD) over the second port of the SSD; and receiving the data from the second storage device at the SSD over the second port of the SSD includes receiving the data from the KV-SSD at the SSD over the second port of the SSD.

(179) Statement 66. An embodiment of the inventive concept includes the method according to statement 65, further comprising generating the KV read request from the first read request.

(180) Statement 67. An embodiment of the inventive concept includes an article, comprising a non-transitory storage medium, the non-transitory storage medium having stored thereon instructions that, when executed by a machine, result in: receiving a first write request for a data from a host at a Solid State Drive (SSD) over a first port of the SSD; sending a second write request for the data from the SSD to a second storage device over a second port of the SSD; and receiving a write result from the second storage device at the SSD over the second port of the SSD; and sending the write result from the SSD to the host over the first port of the SSD.

(181) Statement 68. An embodiment of the inventive concept includes the article according to statement 67, wherein receiving a first write request for a data from a host at a Solid State Drive (SSD) over a first port of the SSD includes receiving the first write request for the data from the host at the SSD via a first queue over the first port of the SSD; and sending the write result from the SSD to the host over the first port of the SSD includes sending the write result from the SSD to the host via the first queue over the first port of the SSD.

(182) Statement 69. An embodiment of the inventive concept includes the article according to statement 68, wherein the first queue includes queue pair including a first submission queue and a

first completion queue.

(183) Statement 70. An embodiment of the inventive concept includes the article according to statement 67, wherein: sending a second write request for the data from the SSD to a second storage device over a second port of the SSD includes sending the second write request for the data from the SSD to the second storage device via a second queue over the second port of the SSD; and receiving a write result from the second storage device at the SSD over the second port of the SSD includes receiving the write result from the second storage device at the SSD via the second queue over the second port of the SSD.

(184) Statement 71. An embodiment of the inventive concept includes the article according to statement 70, wherein the second queue includes queue pair including a second submission queue and a second completion queue.

(185) Statement 72. An embodiment of the inventive concept includes the article according to statement 67, wherein sending a second write request for the data from the SSD to a second storage device over a second port of the SSD includes sending the second write request for the data from the SSD to a plurality of second storage devices over the second port of the SSD.

(186) Statement 73. An embodiment of the inventive concept includes the article according to statement 72, wherein sending the second write request for the data from the SSD to a plurality of second storage devices over the second port of the SSD includes: disassembling the data into at least two portions of the data; and sending a portion of the data from the SSD to one of the plurality of second storage devices over the second port of the SSD.

(187) Statement 74. An embodiment of the inventive concept includes the article according to statement 73, wherein sending the second write request for the data from the SSD to a plurality of second storage devices over the second port of the SSD further includes identifying the one of the plurality of second storage devices using a SPU before sending the portion of the data from the SSD to one of the plurality of second storage devices over the second port of the SSD.

(188) Statement 75. An embodiment of the inventive concept includes the article according to statement 73, wherein disassembling the data into at least two portions of the data includes disassembling the data into at least two portions of the data using a storage processing unit (SPU).

(189) Statement 76. An embodiment of the inventive concept includes the article according to statement 72, wherein receiving the write result from the second storage device at the SSD over the second port of the SSD includes receiving the write result from the plurality of second storage devices at the SSD over the second port of the SSD.

(190) Statement 77. An embodiment of the inventive concept includes the article according to statement 67, the non-transitory storage medium having stored thereon further instructions that, when executed by the machine, result in applying an acceleration function to the data before sending the second write request for the data from the SSD to the second storage device over the second port of the SSD.

(191) Statement 78. An embodiment of the inventive concept includes the article according to statement 77, wherein applying an acceleration function to the data before sending the second write request for the data from the SSD to the second storage device over the second port of the SSD includes applying the acceleration function to the data using a general Data Processing Unit (DPU) before sending the second write request for the data from the SSD to the second storage device over the second port of the SSD.

(192) Statement 79. An embodiment of the inventive concept includes the article according to statement 67, the non-transitory storage medium having stored thereon further instructions that, when executed by the machine, result in storing the data in a flash storage of the SSD.

(193) Statement 80. An embodiment of the inventive concept includes the article according to statement 79, the non-transitory storage medium having stored thereon further instructions that, when executed by the machine, result in sending the write result from the SSD to the host over the first port of the SSD before receiving the write result from the second storage device at the SSD

over the second port of the SSD.

(194) Statement 81. An embodiment of the inventive concept includes the article according to statement 67, wherein: sending a second write request for the data from the SSD to a second storage device over a second port of the SSD sending a Key-Value (KV) write request for the data from the SSD to a Key-Value SSD (KV-SSD) over the second port of the SSD; and receiving a write result from the second storage device at the SSD over the second port of the SSD includes receiving the write result from the KV-SSD at the SSD over the second port of the SSD; and

(195) Statement 82. An embodiment of the inventive concept includes the article according to statement 81, the non-transitory storage medium having stored thereon further instructions that, when executed by the machine, result in generating the KV write request from the first write request.

(196) Statement 83. An embodiment of the inventive concept includes an article, comprising a non-transitory storage medium, the non-transitory storage medium having stored thereon instructions that, when executed by a machine, result in: receiving a first read request for a data from a host at a Solid State Drive (SSD) over a first port of the SSD; determining whether the data is stored in a flash storage of the SSD; based at least in part on the data not being stored in the flash storage of the SSD: sending a second read request for the data from the SSD to a second storage device over a second port of the SSD; and receiving the data from the second storage device at the SSD over the second port of the SSD; and sending the data from the SSD to the host over the first port of the SSD.

(197) Statement 84. An embodiment of the inventive concept includes the article according to statement 83, wherein receiving a first read request for a data from a host at a Solid State Drive (SSD) over a first port of the SSD includes receiving the first read request for the data from the host at the SSD via a first queue over the first port of the SSD; and sending the data from the SSD to the host over the first port of the SSD includes sending the data from the SSD to the host via the first queue over the first port of the SSD.

(198) Statement 85. An embodiment of the inventive concept includes the article according to statement 84, wherein the first queue includes queue pair including a first submission queue and a first completion queue.

(199) Statement 86. An embodiment of the inventive concept includes the article according to statement 83, wherein: sending a second read request for the data from the SSD to a second storage device over a second port of the SSD includes sending the second read request for the data from the SSD to the second storage device via a second queue over the second port of the SSD; and receiving the data from the second storage device at the SSD over the second port of the SSD includes receiving the data from the second storage device at the SSD via the second queue over the second port of the SSD.

(200) Statement 87. An embodiment of the inventive concept includes the article according to statement 86, wherein the second queue includes queue pair including a second submission queue and a second completion queue.

(201) Statement 88. An embodiment of the inventive concept includes the article according to statement 83, wherein sending a second read request for the data from the SSD to a second storage device over a second port of the SSD includes sending the second read request for the data from the SSD to a plurality of second storage devices over the second port of the SSD.

(202) Statement 89. An embodiment of the inventive concept includes the article according to statement 88, wherein receiving the data from the second storage device at the SSD over the second port of the SSD includes receiving the data from the plurality of second storage devices at the SSD over the second port of the SSD.

(203) Statement 90. An embodiment of the inventive concept includes the article according to statement 89, wherein receiving the data from the plurality of second storage devices at the SSD over the second port of the SSD includes: receiving at least two portions of the data from the

plurality of second storage devices at the SSD over the second port of the SSD; and reassembling the data from the at least two portions of the data.

(204) Statement 91. An embodiment of the inventive concept includes the article according to statement 90, wherein reassembling the data from the at least two portions of the data includes reassembling the data from the at least two portions of the data using a storage processing unit (SPU).

(205) Statement 92. An embodiment of the inventive concept includes the article according to statement 88, the non-transitory storage medium having stored thereon further instructions that, when executed by the machine, result in identifying the plurality of second storage devices using a SPU before sending the second read request for the data from the SSD to the plurality of second storage devices over the second port of the SSD.

(206) Statement 93. An embodiment of the inventive concept includes the article according to statement 83, the non-transitory storage medium having stored thereon further instructions that, when executed by the machine, result in applying an acceleration function to the data before sending the data from the SSD to the host over the first port of the SSD.

(207) Statement 94. An embodiment of the inventive concept includes the article according to statement 93, wherein applying an acceleration function to the data before sending the data from the SSD to the host over the first port of the SSD includes applying the acceleration function to the data using a general Data Processing Unit (DPU) before sending the data from the SSD to the host over the first port of the SSD.

(208) Statement 95. An embodiment of the inventive concept includes the article according to statement 83, the non-transitory storage medium having stored thereon further instructions that, when executed by the machine, result in, based at least in part on the data being stored in the flash storage of the SSD, reading the data from the flash storage of the SSD.

(209) Statement 96. An embodiment of the inventive concept includes the article according to statement 83, wherein: sending a second read request for the data from the SSD to a second storage device over a second port of the SSD includes sending a Key-Value (KV) read request for the data from the SSD to a Key-Value SSD (KV-SSD) over the second port of the SSD; and receiving the data from the second storage device at the SSD over the second port of the SSD includes receiving the data from the KV-SSD at the SSD over the second port of the SSD.

(210) Statement 97. An embodiment of the inventive concept includes the article according to statement 96, the non-transitory storage medium having stored thereon further instructions that, when executed by the machine, result in generating the KV read request from the first read request.

(211) Consequently, in view of the wide variety of permutations to the embodiments described herein, this detailed description and accompanying material is intended to be illustrative only, and should not be taken as limiting the scope of the inventive concept. What is claimed as the inventive concept, therefore, is all such modifications as may come within the scope and spirit of the following claims and equivalents thereto.

Claims

1. A device, comprising: storage media for a data; and a controller to receive a first request from a host, generate a second request based at least in part on the first request, send the second request to a storage device, receive a response from the storage device, and send the response to the host, wherein the device includes a cache namespace and a first queue and a second queue for the cache namespace.
2. The device according to claim 1, wherein the device includes: a Peripheral Component Interconnect Express (PCIe) port; or an Ethernet port; or a combination thereof.
3. The device according to claim 1, further comprising a storage-related processing unit (SPU) to manage storage of the data on the storage device.

4. The device according to claim 1, further comprising a general data processing unit (DPU).
 5. The device according to claim 1, wherein the cache namespace includes: a cache manager to manage a third request received via the first queue and to issue a fourth request via the second queue; and a persistent memory to store the data based on the third request and the fourth request.
 6. A device according to claim 1, wherein: the first request includes a read request; the response from the storage device includes the data; and the controller is further configured to store the data in the storage media.
 7. A device according to claim 1, wherein: the response includes the data; and the controller is further configured to store the data in the storage media.
 8. The device according to claim 4, wherein the DPU is configured to execute instructions of an application.
 9. The device according to claim 5, wherein the cache namespace further includes a processing unit (PU) buffer for data assembly and disassembly.
 10. A device according to claim 7, wherein the controller is further configured to receive a third request from the host and to send a second response based at least in part on the data in the storage media.
 11. A method comprising: receiving a first I/O request for a data from a host at a device using a first queue associated with a cache namespace in the device; generating a second I/O request based at least in part on the first I/O request; sending the second I/O request for the data from the device to a storage device using a second queue associated with the cache namespace in the device; receiving a response from the storage device at the device; and sending the response from the device to the host.
 12. The method according to claim 11, wherein sending the second I/O request for the data from the device to the storage device includes sending the second I/O request for the data from the device to the storage device and a second storage device.
 13. The method according to claim 11, wherein: sending the second I/O request for the data from the device to the storage device includes sending a Key-Value (KV) I/O request for the data from the device to a KV storage device; and receiving the response from the storage device at the device includes receiving the response from the KV storage device at the device.
 14. The method according to claim 11, wherein the first I/O request includes a read request or a write request.
 15. The method according to claim 11, further comprising applying an acceleration function to the data based at least in part on receiving the first I/O request for the data from the host at the device.
 16. The method according to claim 12, wherein sending the second I/O request for the data from the device to the storage device and the second storage device includes: disassembling the data into a first portion of the data and a second portion of the data; sending the first portion of the data from the device to the storage device; and sending the second portion of the data from the device to the second storage device.
 17. The method according to claim 12, wherein receiving a response from the storage device at the device includes: receiving a first portion of the data at the device from the storage device; receiving a second portion of the data at the device from the second storage device; and reassembling the data from the first portion of the data and the second portion of the data.
 18. The method according to claim 13, further comprising generating the KV I/O request based at least in part on the first I/O request.
-