

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12399688
Kind Code	B2
Date of Patent	August 26, 2025
Inventor(s)	Sangeneni; Naveen Kumar

Development platform for image processing pipelines that use machine learning with user interface

Abstract

A development platform for implementing a machine learning pipeline on a chip containing multiple hardware compute elements. The development platform includes a user interface, a library of software blocks, and a synthesis engine. The user interface facilitates a user to develop a functional description of the machine learning pipeline. The functional description specifies multiple functional modules, including a machine learning model. The synthesis engine synthesizes the pipeline of functional modules into multiple interconnected executable components of software blocks and generates an implementation package including the executable components and specifying interconnections between the executable components.

Inventors:	Sangeneni; Naveen Kumar (Fremont, CA)
Applicant:	SiMa Technologies, Inc. (San Jose, CA)
Family ID:	1000008777999
Assignee:	SiMa Technologies, Inc. (San Jose, CA)
Appl. No.:	18/132322
Filed:	April 07, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240338179 A1	Oct. 10, 2024

Publication Classification

Int. Cl.: G06F8/20 (20180101); G06F3/04847 (20220101); G06T1/20 (20060101); G06V10/94 (20220101); G06F3/0486 (20130101)

U.S. Cl.:

CPC **G06F8/20** (20130101); **G06F3/04847** (20130101); **G06T1/20** (20130101); **G06V10/945** (20220101); **G06V10/955** (20220101); G06F3/0486 (20130101); G06T2200/24 (20130101); G06T2200/28 (20130101)

Field of Classification Search

CPC: G06F (8/20)

USPC: 717/101

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
10078717	12/2017	Venkataramani	N/A	G06F 30/3315
10360654	12/2018	Maiyuran	N/A	G06F 9/455
10922462	12/2020	Chen	N/A	N/A
2013/0073498	12/2012	Izhikevich et al.	N/A	N/A
2018/0082397	12/2017	Bobrovsky	N/A	G06F 9/544
2018/0113797	12/2017	Breslow et al.	N/A	N/A
2019/0324759	12/2018	Yang et al.	N/A	N/A
2020/0250203	12/2019	Calco et al.	N/A	N/A
2021/0055915	12/2020	Guo	N/A	G06F 8/34
2021/0081836	12/2020	Polleri	N/A	G06F 8/77
2021/0081837	12/2020	Polleri	N/A	G06N 5/022
2021/0273993	12/2020	Shirley et al.	N/A	N/A
2022/0107812	12/2021	Foley	N/A	N/A
2022/0138004	12/2021	Nandakumar	N/A	N/A
2022/0261631	12/2021	Cohen	N/A	G06N 20/00
2022/0283787	12/2021	Tang et al.	N/A	N/A
2023/0080439	12/2022	Saha et al.	N/A	N/A
2023/0083891	12/2022	Achin et al.	N/A	N/A
2023/0161882	12/2022	Kumar	N/A	N/A
2023/0195471	12/2022	Millar	712/22	G06F 9/3877
2023/0305814	12/2022	Sumeet	N/A	G06F 30/343
2024/0193018	12/2023	Situnayake	N/A	G06F 9/5038
2024/0220572	12/2023	Jain	N/A	G06F 17/16
2024/0241651	12/2023	Wang et al.	N/A	N/A
2024/0320036	12/2023	Nandakumar	N/A	G06F 8/34
2024/0338183	12/2023	Sangeneni	N/A	G06F 8/34
2024/0338265	12/2023	Sangeneni	N/A	G06F 8/34
2025/0045492	12/2024	Koiek Akino	N/A	G06F 30/27

OTHER PUBLICATIONS

Richard Lin et al. “Weaving Schematics and Code: Interactive Visual Editing for Hardware Description Languages”—UIST '21, Oct. 10-14, 2021, Virtual Event, USA. cited by examiner
Minutoli, Marco, et al. “Soda: a new synthesis infrastructure for agile hardware design of machine

learning accelerators.” Proceedings of the 39th International Conference on Computer-Aided Design. 2020. cited by examiner

Uliana, David, Peter Athanas, and Krzysztof Kepa. “FPGA-based accelerator development for non-engineers.” 2014 International Conference on ReConFigurable Computing and FPGAs (ReConFig14). IEEE, 2014. cited by examiner

PCT International Search Report and Written Opinion, PCT Application No. PCT/US2024/011655, Apr. 17, 2024, 18 pages. cited by applicant

United States Office Action, U.S. Appl. No. 18/132,319, filed Dec. 18, 2024, 35 pages. cited by applicant

Primary Examiner: Aponte; Francisco J

Attorney, Agent or Firm: Fenwick & West LLP

Background/Summary

BACKGROUND

1. Technical Field

(1) This disclosure relates generally to the implementation of machine learning models on hardware, and more particularly to synthesizing hardware agnostic functional descriptions into a pipeline of executable components that are executed on different hardware compute elements.

2. Description of Related Art

(2) Machine learning (ML) is one of the most powerful recent trends in technology. In machine learning, a model is developed to perform a certain task. The model, which will be referred to as a machine learning network or machine learning model, is trained and deployed in order to carry out that task. For example, a model may be developed to recognize the presence of objects within images captured by a set of cameras. Once the model is deployed, images captured by the cameras are input to the model, which then outputs whether or to what confidence level objects are present within the images.

(3) Image processing pipelines that include machine learning networks may be implemented on different types of hardware, including on chips in edge devices. However, every chip vendor may have their own proprietary hardware with its own compiler. When engineers are faced with a new application, it may take a long time for the engineers to develop the pipeline for the application. Existing development platforms do not provide a way for engineers to easily and quickly realize their solutions in a prototype format.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) Embodiments of the disclosure have other advantages and features which will be more readily apparent from the following detailed description and the appended claims, when taken in conjunction with the example embodiments in the accompanying drawings, in which:
- (2) FIG. 1 is a block diagram of an example system with a vision development platform (VDP) for use with developing a system on chip (SoC).
- (3) FIG. 2 is a block diagram showing an example of a catalog of functional modules, and its corresponding software and graphics libraries.
- (4) FIG. 3 is a block diagram of an example device that includes an SoC with a pipeline manager.

(5) FIG. 4 is a block diagram of an example pipeline manager on an SoC.

(6) FIG. 5A is an example graphical user interface (GUI) of the VDP for users to create a new project.

(7) FIG. 5B is an example GUI of the VDP in canvas view, where users can drag and drop functional modules, including models and apps, from a catalog area onto a canvas area.

(8) FIG. 5C is an example GUI of the VDP in canvas view, where users can link functional modules on a canvas area with arrows.

(9) FIG. 5D is an example GUI of the VDP in canvas view, where users can adjust memory throttling for hardware compute elements.

(10) FIG. 5E is an example GUI of the VDP in canvas view, that allows users to train their own models and import the trained models into the catalog of functional modules.

(11) FIG. 5F is an example GUI of the VDP in source code view, where users can view and edit source code generated by the VDP.

(12) FIG. 5G is an example GUI of the VDP in pipeline view, where users can view a pipeline of interconnected executable components and their corresponding hardware compute elements.

(13) FIG. 5H is an example GUI of the VDP in executable code view, where users can view and install an implementation package onto the SoC.

DETAILED DESCRIPTION OF THE PREFERRED EMBODIMENTS

(14) The figures and the following description relate to preferred embodiments by way of illustration only. It should be noted that from the following discussion, alternative embodiments of the structures and methods disclosed herein will be readily recognized as viable alternatives that may be employed without departing from the principles of what is claimed.

(15) Machine learning networks (MLNs) are commonly implemented in computing facilities with access to significant resources, such as in the cloud or on server clusters. However, the sources of input to ML networks may be located remotely from these large computing facilities. For example, cameras and other types of sensors may be edge devices. Example applications for edge devices include automotive and other forms of transportation including autonomous transportation, agricultural, industrial, robotics, drones, surveillance and security, smart environments including smart cities, medical and personalized health. Example tasks include computer vision, image analysis, image understanding, classification and pattern recognition tasks. For edge devices, it may be desirable to perform certain tasks in real-time. In addition to memory and other programmable processors, an edge device may also include sensors, such as cameras including both still image and video cameras, microphones, temperature sensors, pressure sensors and other types of sensors.

(16) The sensors may capture samples that are used as inputs to a computing pipeline implemented on the edge device. Thus, it would be beneficial if MLNs could also be implemented in edge devices since computing pipelines may include MLNs as one (or more) stages in the pipeline. A machine learning accelerator (MLA) is described herein that may be built into an SoC (system on chip) of an edge device. Additional examples of MLAs and corresponding compilers are described in U.S. Pat. No. 11,321,607, entitled “Machine Learning Network Implemented by Statically Scheduled Instructions, with Compiler,” granted on May 3, 2022.

(17) Different SoCs may have different hardware compute elements, such as MLAs but also including other types of processors. In order to implement a computing pipeline using these hardware elements, engineers must decide which functions should be performed by which hardware compute elements, they must then develop the corresponding software programs including passing data between the different hardware elements, and then they must deploy the entire package on the SoC. This can be a complex task and there can be a long learning curve for the engineers to develop their applications and visualize a proof of concept. Existing development platforms do not provide a way for engineers to easily realize their solutions in a prototype format.

(18) The principles described herein address the above-described problem by providing a development platform (running on a computer system) that allows users to build their pipeline

using a graphical user interface (GUI) without having to write significant amounts of code. The following examples are based on image processing pipelines (including video processing), so the platform is referred to as a “vision development platform” (VDP), but similar development platforms may also be developed for other types of computing pipelines.

(19) The VDP can provide a catalog of functional modules from which the user can assemble their pipeline. Examples include ML models, sensor modules, processing modules, networking modules, applications, and plugins. Functional modules can include modules from open source repositories or other third party sources. The VDP can also suggest networks of functional modules based on desired applications.

(20) The VDP can also check the pipelines formed by users. In some embodiments, the VDP generates modifiable JSON files used to run the pipeline on a chip, and compiles and generates packages for binaries, applications and/or JavaScript object notation (JSON) files. In some embodiments, the VDP also provides build time and run time statistics of how the pipeline will perform on the chip. In some embodiments, the VDP is also able to remotely manage devices running pipelines for users and/or build analytics for the users.

(21) For example, a given chip includes a plurality of hardware compute elements, one of which is an MLA. A user wants to use the chip to implement an ML pipeline, which is a computing pipeline that uses a machine learning model. The ML pipeline may be an image processing pipeline, for example. The user can use the VDP and its catalog of functional modules to quickly and easily design a pipeline for execution on the SoC without having to know the details of the hardware compute elements on the SoC.

(22) The VDP may include a catalog of functional modules, and a library of corresponding software blocks that implement the functional modules. Each of the software blocks corresponds to an atomic functional stage of a functional module that is to be executed by a hardware compute element. Some functional modules may include a single software block, so that the entire functional module is executed by a single hardware compute element. Other functional modules may include multiple software blocks, so that different parts of the functional module are executed by different hardware compute elements.

(23) A user can enter a functional description of a computing pipeline that specifies the functional modules that form the pipeline, e.g., ML models, sensor device functional modules, input/output (I/O) device functional modules, etc. The functional description may be hardware agnostic, i.e., the user does not need to specify which hardware compute element is to perform which part of the functional module or pipeline. The entering of the functional description of the computing pipeline can be performed via a GUI or descriptive language. Based on the user input, the VDP accesses the software library, retrieves the software blocks corresponding to the functional modules in the pipeline, and compiles the software blocks into executable components for corresponding hardware compute elements.

(24) The VDP then generates an implementation package that includes the executable components and specifies the interconnections between them. In some embodiments, the interconnections are described in JSON files. The implementation package includes the executable components and the JSON files. The implementation package can then be deployed onto the SoC. The SoC includes a pipeline manager that parses the implementation package and distributes the different executable components to different hardware compute elements for execution in a proper sequence. As such, a user is able to use functional descriptions to develop application projects that can be executed on the SoC without having to learn the specifics of the proprietary hardware of the SoC. For example, the user does not need to know how an ML model is partitioned into software blocks or which hardware compute element on the SoC executes each of the software blocks.

(25) FIG. 1 is a high-level diagram of such a system. The system includes a VDP **110** for use with a chip **180**. The VDP **110** includes a synthesis engine **120** that receives the functional description **140** of an ML pipeline from a user **160**. The functional description **140** specifies the functional modules

(from catalog **190**) that form the ML pipeline. At least one of the functional modules includes an ML model. The synthesis engine **120** synthesizes the functional description **140** into interconnected executable components (part of implementation package **170** in FIG. **1**), each of which is to be executed by a particular hardware compute element of the chip **180**.

(26) In some embodiments, synthesis engine **120** has access to catalog **190** of the functional modules, a software library **197**, and a hardware compute element listing **199**. The functional module catalog **190** includes names and/or descriptions of multiple functional modules. The functional modules may include ML models **192**, sensor device functional modules **194**, I/O device functional modules **196**, etc. The software library **197** includes software blocks used to implement the functional modules in the catalog **190**. In some embodiments, the software blocks include source code files written in one or more particular computer-programming languages, such as C, C++, Java, etc. The hardware compute element listing **199** includes descriptions of multiple hardware compute elements that are implemented in the chip **180**. Such hardware compute elements may include various application processing units (APU) MLAs, computer vision units (CVUs), and other processors or compute elements.

(27) The synthesis engine **120** receives the functional description **140** of a computing pipeline, which includes multiple functional modules and their interconnections. The synthesis engine **120** accesses the functional module catalog **190** and the software library **197** to retrieve the software blocks corresponding to the functional modules. As discussed above, certain functional modules may include multiple functional stages executed on different hardware compute elements. Such a functional module corresponds to multiple software blocks, each of which is compiled separately to generate a separate executable component. The synthesis engine **120** maps each of the executable components to a particular hardware compute element implemented in the chip **180**. The executable components and their interconnections are then packaged into an implementation package **170** and deployed onto the chip **180**.

(28) For example, the computing pipeline may include an ML model. The ML model may include three interconnected software blocks, one of which is to be executed by an MLA of the chip, and the other two are to be executed by an application processing unit (APU) of the chip. The synthesis engine **120** also connects the executable components corresponding to different functional modules into a pipeline of interconnected executable components. The interconnections among the executable components are written in a particular format and stored in files, such as in JSON file(s). The files and the executable components are packaged into an implementation package **170** and deployed onto the chip **180**.

(29) In some embodiments, the synthesis engine **120** includes one or more frontend modules **122**, a compiler module **127**, and one or more backend modules **128**. The front end modules **122** for ML models include pruning, compression, and quantization modules **124** and a partition module **126**. The pruning module **124** removes parts of the ML model that do not contribute significantly to the overall results. The quantization module **124** reduces the resolution of calculated values. Because ML models contain a large amount of data, the compression module **124** may be used to reduce data transfer bandwidths.

(30) As discussed above, certain functional modules may include multiple stages, which are mapped to software blocks that are executed on hardware compute elements of the chip **180**. The partition module **126** partitions certain ML models into multiple stages. In some embodiments, the partition and mapping of the different stages may be based on specializations of each hardware compute element implemented on the chip **180**. For example, an ML model may be partitioned into a tensor multiplication block and a nonlinear operator block. The tensor multiplication block may be mapped to an MLA for execution, and the nonlinear operator block may be mapped to an APU for execution.

(31) The compiler module **127** compiles the software blocks for the different functional modules into executable components. Each of the executable components is executable on a particular

hardware compute element. The backend module **128** performs operations after the compilation of the source code. For example, the backend module **128** may include a pipeline generator that links the executable components in a particular sequence and generates the implementation package **170** containing the executable components and their interconnections. The synthesis engine **120** may also include additional modules or functions **129**.

(32) In some embodiments, VDP **110** provides a graphical user interface (GUI) **150** that a user can use to provide specifications of the chip **180** and the functional description **140** of a computing pipeline. The specifications of the chip **180** include one or more hardware compute elements implemented on the chip **180**. In some embodiments, VDP **110** has access to a graphics library **198** that stores graphics representing the functional modules. The VDP **110** allows the user to visualize the pipeline in the GUI, using the graphics corresponding to the functional modules of the pipeline.

(33) In some embodiments, the GUI **150** displays the catalog **190** of functional modules. A user can select one or more functional modules from the displayed catalog. In some embodiments, the GUI **150** includes a canvas view that allows the user **160** to drag and drop functional modules from the catalog onto a canvas area. When a functional description is dragged onto the canvas area, VDP **110** accesses the graphics library **198** to retrieve a graphic corresponding to the functional module and display the graphic on the canvas. The user can then link the graphics with connectors (e.g., lines and arrows) to indicate connections between the functional modules. In some embodiments, GUI **150** includes a code view that allows the user to view and edit the corresponding software code.

(34) An example of a relationship between a functional module, and the corresponding graphics and software blocks are further discussed below with respect to FIG. 2. FIG. 2 illustrates an example of mapping a functional module **210** to software blocks **202**, **204**, **206** and a graphic **220**. As illustrated in FIG. 2, the functional module is “CenterNet,” which is a pre-trained ML network. CenterNet includes three functional stages CenterNet_1, CenterNet_2 and CenterNet_3, each of which is to be executed on a different hardware compute element in a particular sequence. As such, each functional stage corresponds to a separate software block **202**, **204**, **206** that is compiled and executed on the corresponding hardware compute element. The “CenterNet” functional module is represented by graphic **220**, which includes three sub-graphics **222**, **224**, **226** representing the three functional stages.

(35) A user does not need to understand how many functional stages CenterNet has, and which hardware compute element is to execute which functional stage. Instead, the user selects the hardware compute elements that are implemented on the chip (e.g., an APU, an MLA, a CVU), and inputs the functional description of the functional module, i.e., “CenterNet.” Based on the user input, VDP **110** automatically partitions the CenterNet into three functional stages, and maps the three functional stages to different compute elements. In this example, CenterNet_1 is implemented by software block **202** executing on the APU, CenterNet_2 is implemented by software block **204** executing on the MLA, and CenterNet_3 is implemented by software block **206** executing on the APU.

(36) In some embodiments, VDP **110** assembles the software blocks **202**, **204**, **206** into a set of source code files for the user's project. Similarly, other functional modules are mapped to their corresponding software blocks and graphics. A user can input a functional description of an ML pipeline by selecting and interconnecting multiple functional modules from the catalog **190**. Based on the user input, VDP **110** generates source code and then executable components based on the functional description of the ML pipeline. As such, the user can create complex ML pipelines without writing significant amounts of code.

(37) Returning back to FIG. 1, in some embodiments, VDP **110** also includes a key performance indicator (KPI) calculator **130** that calculates values that measure the performance of the ML pipeline corresponding to the functional description **140**. These KPI values can be displayed in the GUI **150**. Examples of KPIs include frames per second (FPS), power consumption, memory

utilization, and processor utilization.

(38) In some embodiments, VDP **110** is installed on a client device of a user, and the user can deploy the implementation package onto a chip **180** (also referred to as a “target chip”) by connecting the chip **180** to the client device, e.g., via wired or wireless communications. In some embodiments, VDP **110** is a cloud system that is physically connected to various chips. A user can remotely access VDP **110** and cause the VDP to deploy the implementation package onto a target chip that is connected to VDP. In some embodiments, VDP **110** includes or is coupled to an emulator or simulator that emulates or simulates various chips. The implementation package may be deployed onto an emulation or simulation of the target chip.

(39) As briefly discussed above, once the chip receives the implementation package, a pipeline manager of the chip parses the implementation package and causes the different executable components to be executed by different hardware compute elements in proper order. FIG. **3** is a block diagram of an example device **300** with a system-on-chip (SoC) **302**, which is an example of chip **180** of FIG. **1**. The SoC **302** includes an MLA **370**. Other components may be included on the same die as the MLA. This example includes the following additional blocks: application processing unit (APU) **310** (e.g., general purpose CPU running applications), computer vision unit (CVU) **312** (or other types of application-specific processors), safety **314**, security **316**, additional SRAM (memory) **320** and input/output (I/O) circuitry **324**. It also includes a network **330** for communication between the different components. The connections to the external world include camera inputs **340** for the CVU **312**, ports for debugging **342** and configuration **344**, a connection **346** to external memory (e.g., DRAM), chip-to-chip connections **348**, and network connections **350** (e.g., Ethernet and PCIe).

(40) A pipeline manager **322** is installed on the SoC **302** and executable by APU **310**. The pipeline manager **322** interprets the implementation package **170** received from the VDP **110**. As discussed above with respect to FIG. **1**, the implementation package **170** includes multiple executable components and specification of the interconnections between components. The pipeline manager **322** parses the implementation package **170** to extract the interconnected executable components, and distribute them to their respective hardware compute elements for execution, such as on CVU **312**, MLA **370**, etc.

(41) FIG. **4** is a block diagram of an example pipeline manager **322** (e.g., GStreamer) that manages a pipeline **400** of executable components **401-408**. In FIG. **4**, each block includes a hardware compute element and a cycle number, indicating that during the cycle, the executable component is to be executed on the hardware compute element. For example, executable component **401** is executable on CVU **312**; executable component **402** is executable on APU **310**; executable component **403** is executable on MLA **370**, and so on and so forth.

(42) The pipeline manager **322** manages the timing and the location of execution of the executable components **401-408** based on information in the implementation package **170**. In this example, executable components **401-403** are executed starting in cycle **0**. Executable components **404-405** are executed starting in cycle **1**. Executable components **406-407** are executed starting in cycle **2**. Executable component **408** is executed starting in cycle **3**. The components **401-408** are connected in a pipeline as shown.

(43) In some embodiments, after the executable component **408** is executed, a new round of operations may be performed, starting from block **401** again as indicated by the dashed arrow. For example, the pipeline **400** may perform object recognition to identify an object from a video stream. The pipeline **400** may be tasked to constantly monitor the frames of images in the video stream to identify the object. After a first frame of image is processed, the pipeline is executed again to process a second frame of image.

(44) Referring back to FIG. **3**, in addition to memory and other programmable processors, the edge device **300** may also include sensors, such as cameras (both still image and video cameras), microphones, temperature sensors, pressure sensors, and other types of sensors. The sensors may

capture samples that are used as inputs to a pipeline within the edge device. For example, image samples may be input to the CVU **312**, which performs initial operations such as edge detection and enhancement, contrast enhancement, motion detection, and optical flow. These may be earlier functional modules in the pipeline. Raw and/or processed images may then be input to the MLA **370** for analysis by the ML network. The MLA **370** may also receive other inputs, such as metadata from other sources and data from other sensors. The APU **310** may also perform various functions in the overall pipeline and serve as a master controller that coordinates the operation of the MLA **370** and the other hardware compute elements in the pipeline.

(45) Example applications for edge device **300** include automotive and other forms of transportation including autonomous transportation, agricultural, industrial, robotics, drones, surveillance and security, smart environments including smart cities, medical and personalized health. Example tasks include computer vision, image analysis, image understanding, speech recognition, audio analysis, audio understanding, natural language processing, classification and pattern recognition tasks.

(46) Traditionally, a user would have to understand details about various software functions and various hardware compute elements on the SoC **302**, so that the user can write source code for the software functions that are to be executed on different hardware compute elements. There is a steep learning curve for even experienced engineers to be able to grasp the nuances of each SoC.

(47) VDP **110** solves this problem by providing an interface in which a user provides functional descriptions of different processes (i.e., functional modules). The VDP synthesizes the functional modules of the pipeline into a plurality of interconnected executable components, which can be deployed onto the SoC. As such, users do not have to understand the details of various software functions and the different hardware compute elements. The functional descriptions may be entered via text format, such as JSON code, or any descriptive language. Alternatively, or in addition, the functional descriptions may be entered via drag and drop of graphics representing different functions onto a canvas area of a GUI.

(48) FIG. 5A illustrates an example graphical user interface (GUI) of VDP **110** for a user to create a new project or application. The user can enter a project name **510**, select a board type **511**, select different hardware compute elements (e.g., APU, CVU, MLA, and other compute elements or processors) **512** in a chip, select operating systems **513** (e.g., Linux, RTOS, etc.), and select cores **514** for an SoC. The user can also choose to install additional packages and libraries **515** onto the SoC, such as an additional operating system. The user may also create new projects based on recent projects or templates **516**. As described in more detail below, the user can also specify different pipelines to be implemented on the SoC. When the new project is created, VDP **110** generates a set of source code from a software library. The library may include a set of folders containing a set of source code files, which may be written in various programming languages, such as C, C++, C#, java, JavaScript, JSON, etc.

(49) VDP **110** can also generate different views of the project, such as canvas view, source code view, pipeline view, or executable code view. The canvas view is a GUI that looks like a canvas, and a user can generate a functional description of an ML pipeline by dragging and dropping different functional modules onto the canvas, and linking the functional modules on the canvas. When an additional functional module is dragged onto the canvas, or an additional link is created between two different functional modules, VDP **110** modifies the set of source code, causing the source code to include the corresponding software blocks and their interconnections.

(50) The code view is a GUI that looks like a code editor, and a user can review and edit the set of source code generated by VDP **110**. The pipeline view is a GUI that shows a pipeline of interconnected software blocks corresponding to the project. The pipeline view may be generated after the source code is compiled into executable code including multiple executable components, each of which is executable on a particular hardware compute element. The compiled code may be viewed and deployed onto an SoC in the executable code view. The interconnections among the

multiple executable components may be presented in JSON format. The JSON code may also be viewed and edited via the executable code view.

(51) FIGS. 5B-5E illustrate example GUIs of VDP **110** in the canvas view. FIG. 5B illustrates a GUI **500B**. On the right side of the GUI **500B**, there is a models/apps catalog **520** showing a list of ML models and other functional modules. In FIG. 5B, the user has selected a 2D convolutional ML network, called CenterNet **521**, as indicated by the cross-hatch. The user may click “CenterNet” in the catalog to select it, or drag it onto the canvas area **522**. Once the CenterNet is selected or placed on the canvas, the VDP **110** updates the canvas area to show a graphic **523** representing the CenterNet functional module. The CenterNet module includes three software blocks, each of which is to be performed by a particular hardware compute element. The graphic **523** includes graphics representing the three software blocks and indicating the corresponding hardware compute element. In this example, the first and third functional stages are executed by the APU, while the second functional stage is executed by the MLA. The VDP **110** also updates the set of source code to include the software blocks corresponding to CenterNet.

(52) Note, the user does not need to know how many functional stages (software blocks) the CenterNet has, or which hardware compute element of the SoC executes which functional stage. In response to the user's drag and drop of the CenterNet into the canvas area, VDP **110** automatically updates the source code to include the software blocks corresponding to CenterNet, which includes the three blocks. VDP **110** maps each of the three blocks to a particular hardware compute element implemented in the SoC. The hardware compute elements of the SoC may be automatically set by VDP **110** or selected by the user. In some embodiments, VDP **110** may consider different hardware compute elements for each block. The VDP **110** represents the CenterNet on the canvas area using a graphic that includes the three functional stages and their corresponding hardware compute elements.

(53) In some embodiments, the VDP **110** also computes various key performance indicators (KPIs) based on the generated pipeline. As shown in area **524** of FIG. 5B, the GUI shows frame per second (FPS), power utilization, memory utilization, each hardware compute element's utilization.

(54) FIG. 5C illustrates a GUI **500C** showing a list **530** of functional modules for sensors and ethernet devices that may be used in the project. A user can select one or more of these sensors or ethernet devices as a data source (e.g., drag and drop a sensor or ethernet device into the canvas area), and connect the functional modules in the canvas area (e.g., by linking them with arrows). Again, a user does not need to know how many functional stages are in each functional module, or which hardware compute element executes the corresponding software. VDP **110** automatically modifies the source code to reflect the added sensors or other devices and recomputes KPIs, while updating the graphic in the canvas area.

(55) In some embodiments, VDP **110** detects incorrect connections made by users. For example, when a user links two functional modules that are not supposed to be linked together, or the linking direction is incorrect, VDP **110** may generate a warning to alert a user. For example, when a user links an ML model output to a sensor block input, VDP **110** may generate an alert, suggesting that the user changes the arrow direction to link the sensor block output to the ML model input.

(56) FIG. 5D illustrates a GUI **500D**, in which a user is able to adjust memory throttling for each hardware compute element using slider **540**. Memory throttling restricts read and/or write traffic to main memory as a means of controlling power consumption. The KPIs are updated based on the adjusted memory throttling. Again, VDP **110** would update the source code to reflect the adjustment of the memory throttling. Power consumption is an important concern for applications deployed on edge devices. Thus, it is advantageous for a user to be able to see correlations between power consumption and performance during the design stage, and adjust memory throttling of hardware compute elements accordingly.

(57) FIG. 5E illustrates a GUI **500E**, in which a user is provided an option to train their own models. In some embodiments, the training can be performed by an AI server that is associated

with VDP **110**. Alternatively, the training can be performed on a client device of the user. The user can configure various hyperparameters for the training, as shown by **550**. Once the model is trained, the model can be added to the library or catalog of VDP **110**.

(58) FIG. 5F illustrates an example GUI **500F** of VDP in the source code view. As shown in FIG. 5F, the source code is written in C programming language. Each of the functional modules, models, sensors, and/or devices includes software blocks, which are software components that can be added or linked to the C source code framework. Each of the software block that is added or linked to the project is also referred to as a “plugin,” which corresponds to a JSON file. The JSON file describes the relationships between the software block and other source code associated with the project. A list of source code files and plugins are shown on a left panel **560** of the GUI. A source code viewer or editor is shown on the right side **561** of the GUI. When a user selects one of the source code files on the left panel **560**, the selected source code file is displayed on the viewer/editor side **561**. It is advantageous for expert users to be able to view and modify source code, although no modification or coding is required for a user to create and execute an application.

(59) After a user finishes their design of a project, VDP **110** compiles the source code into executable components, and packages the executable components into an implementation package based on their interconnections. In some embodiments, once the source code is compiled, VDP **110** can generate a pipeline view of the application, showing the interconnections of each executable component.

(60) FIG. 5G illustrates an example GUI **500G** of VDP in the pipeline view. The GUI **500G** shows a pipeline of multiple interconnected blocks, each of which corresponds to an executable component, and a hardware compute element. The executable component is executed on the hardware compute element. For example, block **570** executes on the APU, block **571** executes on the CVU, block **572** executes on the MLA, and so on and so forth. Further, the arrows linking different executable components represent the interconnections between components. For example, data flows from block **570** to block **571** to block **572**, etc. The tee block **574** splits the pipeline into two branches. The top branch includes blocks **575** and **576**. The bottom branch includes block **577**. Both branches flow into block **578**.

(61) FIG. 5H illustrates an example GUI **500H** of VDP in the executable code view. This GUI **500H** includes a left panel **580** that shows a list of files that are related to the implementation package, and a viewer or editor **581** on the right. A user can click “install implementation package” link **582** to cause the implementation package to be installed or deployed onto the SoC. In some embodiments, the implementation package includes multiple executable components and a JSON file that describes relationships among the multiple executable components. When the user selects the implementation package on the left panel, the JSON file is shown in the viewer or editor. For example, as illustrated in FIG. 5H, the implementation package (in JSON format) includes two executable files, CVU_3.bin and APU_3.bin. CVU_3.bin is executable by a CVU of the SoC, and APU_3.bin is executable by an APU of the SoC. A user can modify the JSON file, while no modification or coding is required to get the implementation package to be deployed onto the SoC.

(62) In some embodiments, VDP **110** also automatically generates documentation for the implementation package. The documentation describes the functions of each executable component and/or their corresponding source code. A user can read the documentation to better understand the functions and interrelations among the different plugins integrated into the project.

(63) Although the detailed description contains many specifics, these should not be construed as limiting the scope of the invention but merely as illustrating different examples. It should be appreciated that the scope of the disclosure includes other embodiments not discussed in detail above. Various other modifications, changes and variations which will be apparent to those skilled in the art may be made in the arrangement, operation and details of the method and apparatus disclosed herein without departing from the spirit and scope as defined in the appended claims.

Therefore, the scope of the invention should be determined by the appended claims and their legal equivalents.

Claims

1. A development platform for implementing a machine learning pipeline on a chip containing a plurality of hardware compute elements, the development platform comprising: a user interface that facilitates a user to develop a functional description of the machine learning pipeline, wherein the functional description specifies a plurality of functional modules that form the pipeline, and at least one functional module includes a machine learning model; a library of software blocks that implement the functional modules on the hardware compute elements of the chip, wherein the library of software blocks includes software blocks for implementing a plurality of machine learning models, a plurality of sensor plugins, and a plurality of ethernet device plugins; and a synthesis engine configured to: synthesize the pipeline of functional modules into a plurality of interconnected executable components of the software blocks, wherein the machine learning pipeline is synthesized into executable components that execute on at least two different hardware compute elements of the chip; and generate an implementation package comprising the executable components and specifying interconnections between the executable components.
2. The development platform of claim 1 further comprising: a key performance indicator (KPI) calculator that computes a KPI based on the interconnected executable components the KPI displayed in the user interface.
3. The development platform of claim 2 wherein the KPI includes at least one of frames per second (FPS), power consumption, memory utilization, and processor utilization.
4. The development platform of claim 2 wherein the KPI includes KPIs for each of the hardware compute elements.
5. The development platform of claim 2 wherein, responsive to the user's modification of the pipeline of functional modules, the KPI calculator updates the KPI.
6. The development platform of claim 1 wherein the user interface includes a graphical user interface (GUI) that displays a catalog of functional modules and allows the user to select functional modules from the catalog.
7. The development platform of claim 6 wherein the GUI further includes a canvas area; and, responsive to the user dragging functional modules from the catalog to the canvas area, the GUI displays the dragged functional modules in the canvas area.
8. The development platform of claim 6 wherein one of the selected functional modules includes a plurality of submodules that execute on corresponding hardware compute elements, and the GUI displays the plurality of submodules and the corresponding hardware compute elements.
9. The development platform of claim 6 wherein the GUI displays one or more parameters of a hardware compute element, and allows the user to modify the one or more parameters of the hardware compute element.
10. The development platform of claim 1 wherein the synthesis engine is further configured to generate source code files for software blocks implementing the pipeline of functional modules, and compile the source code files into the plurality of interconnected executable components.
11. The development platform of claim 1 wherein the synthesis engine is further configured to: determine whether a connection between two functional modules is correct, and responsive to determining an incorrect connection, generate an error alert.
12. The development platform of claim 1 further comprising a model training engine configured to: train a custom machine learning model based on user input, and include software blocks for implementing the custom machine learning model in the library.
13. The development platform of claim 1 wherein the library of software blocks includes proprietary software blocks, open source software blocks, and user-generated software blocks.

14. The development platform of claim 1 further comprising a library of different chips comprising different processors, and a user is allowed to select a target chip from the library of chips, and the synthesis engine generates the executable components based on the processors on the target chip.
15. The development platform of claim 1 wherein the library of software blocks comprises software blocks for implementing a plurality of machine learning models.
16. The development platform of claim 1 further comprising: a library of reference pipelines.
17. The development platform of claim 1 wherein the development platform supports a concurrent design of multiple different chips and multiple different versions of chips.
18. The development platform of claim 1 wherein the development platform is a cloud-based service.
19. The development platform of claim 1 wherein the development platform is connected to the chip, and configured to receive a KPI of the chip when the executable components are executed by the hardware compute elements of the chip.
20. A development platform for implementing a machine learning pipeline on a chip containing a plurality of hardware compute elements, the development platform comprising: a user interface that facilitates a user to develop a functional description of the machine learning pipeline, wherein the functional description specifies a plurality of functional modules that form the pipeline, and at least one functional module includes a machine learning model; a library of software blocks that implement the functional modules on the hardware compute elements of the chip; a library of different chips comprising different processors; and a synthesis engine configured to: synthesize the pipeline of functional modules into a plurality of interconnected executable components of the software blocks, wherein the machine learning pipeline is synthesized into executable components that execute on at least two different hardware compute elements of the chip; and generate an implementation package comprising the executable components and specifying interconnections between the executable components, wherein a user is allowed to select a target chip from the library of chips, and the synthesis engine generates the executable components based on the processors on the target chip.
21. A development platform for implementing a machine learning pipeline on a chip containing a plurality of hardware compute elements, the development platform comprising: a user interface that facilitates a user to develop a functional description of the machine learning pipeline, wherein the functional description specifies a plurality of functional modules that form the pipeline, and at least one functional module includes a machine learning model; a library of software blocks that implement the functional modules on the hardware compute elements of the chip, wherein the development platform supports a concurrent design of multiple different chips and multiple different versions of chips; and a synthesis engine configured to: synthesize the pipeline of functional modules into a plurality of interconnected executable components of the software blocks, wherein the machine learning pipeline is synthesized into executable components that execute on at least two different hardware compute elements of the chip; and generate an implementation package comprising the executable components and specifying interconnections between the executable components.
-