

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12399703
Kind Code	B2
Date of Patent	August 26, 2025
Inventor(s)	Suresh; Prashanth et al.

Parallel multi-rack dataplane upgrades

Abstract

Examples described herein include systems and methods for parallel upgrades of multiple hosts within an on-premises, multi-rack software-defined data center (“SDDC”). Upgrade instructions can be provided to a point-of-presence (“POP”) virtual machine (“VM”) of the SDDC, which can call an application programming interface (“API”) of a pod service executing on a control plane VM in the SDDC. For example, the pod service can generate a set of instructions for performing cluster upgrades in parallel. In one example, the pod service provides the instructions to a lifecycle management service executing on a VM of the SDDC, such as the control plane VM. In some examples, the lifecycle management service can add a spare host to a cluster of hosts in order to begin upgrading one or more hosts in that cluster. This can be performed in parallel, using the spare hosts from different racks simultaneously.

Inventors: Suresh; Prashanth (San Jose, CA), Munot; Sushil Shripal (Santa Clara, CA), Vellore; Shyama Sankar (Mountain View, CA), Rajurkar; Rucha (San Jose, CA)

Applicant: VMware LLC (Palo Alto, CA)

Family ID: 1000008778378

Assignee: VMware LLC (Palo Alto, CA)

Appl. No.: 18/315899

Filed: May 11, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240378044 A1	Nov. 14, 2024

Publication Classification

Int. Cl.: G06F9/48 (20060101); G06F8/65 (20180101); G06F8/656 (20180101); G06F9/50 (20060101); G06F8/61 (20180101); G06F8/71 (20180101); G06F9/455 (20180101)

U.S. Cl.:

CPC G06F8/65 (20130101); G06F8/63 (20130101); G06F8/656 (20180201); G06F8/71 (20130101); G06F9/45558 (20130101); G06F9/4881 (20130101); G06F9/5027 (20130101); G06F9/5077 (20130101)

Field of Classification Search

CPC: G06F (8/65); G06F (8/656); G06F (8/71); G06F (8/63); G06F (9/4881); G06F (9/5027); G06F (9/45558); G06F (9/5077)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2019/0220271	12/2018	Olderdissen	N/A	G06F 9/45558
2020/0201665	12/2019	Panse	N/A	G06F 9/5077
2023/0095814	12/2022	Sarkar	707/641	G06F 11/1451
2023/0289216	12/2022	Agrawal	N/A	G06F 9/4881

OTHER PUBLICATIONS

Vijayaraghavan Soundararajan et al., The Impact of Management Operations on the Virtualized Datacenter, Jun. 19-23, 2010, [Retrieved on Mar. 12, 2025]. Retrieved from the internet: <URL: <https://dl.acm.org/doi/pdf/10.1145/1816038.1816003>> 12 Pages (326-337) (Year: 2010). cited by examiner

Primary Examiner: Rivera; Anibal

Attorney, Agent or Firm: King Intellectual Asset Management

Background/Summary

BACKGROUND

(1) Although cloud computing continues to rapidly gain in popularity and market share, physical datacenter hardware is still required in certain circumstances. For example, a business with highly sensitive data might choose to utilize a physical datacenter at the business location, also known as an “on-premises” datacenter. In another example, a remote edge location such as a cell tower might not have the necessary internet access to rely on a cloud-based datacenter. In these examples and others, a physical, on-premises datacenter is required.

(2) These datacenters can be software-defined data centers (“SDDCs”), which use software to abstract and automate the management of infrastructure resources such as compute, storage, and networking. An SDDC typically consists of several layers of software, including virtualization software to abstract compute resources, software-defined storage to manage storage resources, and software-defined networking to manage network resources. These layers are managed through a centralized management platform.

(3) Physical SDDCs typically include multiple data center racks, also known as server racks. A rack generally refers to a grouping of servers and/or networking equipment that typically reside in a common physical frame. For example, based on a customer's datacenter requirements, they may order multiple racks of servers with which to construct an SDDC.

(4) Various components of an SDDC, such as the data plane components, need to be upgraded over time to maintain performance. Data plane components can be any components within the SDDC architecture that handle the movement and processing of data between applications and infrastructure, performing tasks such as processing and forwarding network traffic, storing and retrieving data, and executing compute tasks. Example SDDC components that are considered part of the data plane can include: compute resources (the virtual machines (“VMs”) or containers that execute applications and services); storage resources (the storage devices that store data and provide data services such as backup, replication, and encryption); network resources (the network devices that handle traffic and provide network services such as load balancing, security, and segmentation); hypervisors or virtualization software (the software that creates and manages VMs and virtualizes compute and storage resources); and automation and orchestration software (the software that automates the deployment, configuration, and management of infrastructure resources and applications).

(5) To upgrade these various components, each host, or server, of the SDDC may need to be upgraded. Generally, each server rack includes one spare host that can be used during upgrades. Data residing on a host needing upgrades can be migrated to the spare host until the upgrade is complete, after which the data can be moved back to the upgraded host. However, today's on-premises SDDCs can only upgrade one host at a time regardless of how many racks or spare hosts are in the SDDC. This means that for a 10-rack datacenter with 10 spare hosts, for example, only one of those spare hosts can be used at any given time for upgrades. This limitation makes upgrades extremely time consuming.

(6) As a result, a need exists for systems and methods for more efficiently upgrading an on-premises, multi-rack SDDC. In particular, a need exists for systems and method for parallel upgrades of multiple hosts within an on-premises, multi-rack SDDC, without impacting customer resources or workloads.

SUMMARY

(7) Examples described herein include systems and methods for parallel upgrades of multiple hosts within an on-premises, multi-rack SDDC. In an example method, upgrade instructions can be provided to a point-of-presence (“POP”) VM. The POP VM can be deployed at a location where network traffic is exchanged between the SDDC and different networks, such as a management service or backend that provides instructions and upgrades to the SDDC. The management service can generate instructions for a data plane upgrade job and send it to the SDDC by way of the POP VM.

(8) The POP VM can call an application programming interface (“API”) of a pod service executing on a management VM (also referred to herein as a control plane VM) in the SDDC. The pod service can include logic for parsing instructions and coordinating specific actions within the SDDC to carry out those instructions. For example, the pod service can generate a set of instructions for performing cluster upgrades in parallel. The instructions can identify multiple clusters, where each cluster includes multiple hosts, and where each cluster is located on a separate data center rack. The pod service can then provide the instructions to the components responsible for carrying them out to perform the upgrades.

(9) In one example, the pod service provides the instructions to a lifecycle management service executing on a VM of the SDDC, such as the management VM. The pod service can invoke APIs of the lifecycle management service to accomplish this. For example, the pod service can generate multiple sets of parallel instructions for performing various upgrades and can invoke multiple APIs of the lifecycle management service in parallel for each of the instruction sets.

(10) In some examples, the lifecycle management service can add a spare host to a cluster of hosts in order to begin upgrading one or more hosts in that cluster. This can be performed in parallel, using the spare hosts from different racks at the same time. For example, a first spare host from a first rack can be added to a cluster executing on the first rack, while a second spare host from a second rack can be added to a cluster executing on the second rack. The lifecycle management service can cause data to be transferred from one or more hosts of each cluster to the spare host added to that cluster. After being backed up using the spare host, the hosts can be taken offline for upgrades.

(11) Upgrades can include upgrading the software, firmware, or both, of the offline hosts. In some examples, the software upgrades are provided by the lifecycle management service directly, such as by providing upgrade files and instructions for applying the upgrade files to multiple offline hosts on different racks. This can cause the offline hosts to upgrade themselves at the same time, using different spare hosts on the different racks.

(12) In some examples, the lifecycle management service can instruct one or more hardware security modules to perform the firmware upgrades. In one example, each cluster includes a firmware manager VM executing in that cluster. Each firmware manager VM can include a hardware security module. To perform parallel firmware upgrades, the lifecycle management service can provide parallel instructions to hardware security modules associated with each cluster to which a spare host has been added. Each hardware security module can perform the firmware upgrades one or more hosts within the same cluster as the firmware manager VM.

(13) The method can also include bringing the offline hosts back online. When all hosts in a cluster have been upgraded, or when the upgrade instructions have otherwise been completed, the spare host can be removed from that cluster. The process can then repeat for other clusters as appropriate.

(14) The examples summarized above can each be incorporated into a non-transitory, computer-readable medium having instructions that, when executed by a processor associated with a computing device, cause the processor to perform the stages described. Additionally, the example methods summarized above can each be implemented in a system including, for example, a memory storage and a computing device having a processor that executes instructions to carry out the stages described.

(15) Both the foregoing general description and the following detailed description are exemplary and explanatory only and are not restrictive of the examples, as claimed.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1 is a flowchart of an example method for parallel upgrades of multiple hosts within an on-premises, multi-rack SDDC.

(2) FIG. 2 is a sequence diagram of an example method for parallel upgrades of multiple hosts within an on-premises, multi-rack SDDC.

(3) FIG. 3 is a flowchart of an example method for parallel upgrades of multiple hosts within an on-premises, multi-rack SDDC.

(4) FIG. 4 is a flowchart of an example method for parallel upgrades of multiple hosts within an on-premises, multi-rack SDDC.

(5) FIG. 5 is an illustration of an example system for performing the methods described herein.

(6) FIG. 6 is an illustration of an example SDDC for performing the methods described herein.

DESCRIPTION OF THE EXAMPLES

(7) Reference will now be made in detail to the present examples, including examples illustrated in the accompanying drawings. Wherever possible, the same reference numbers will be used throughout the drawings to refer to the same or like parts.

(8) Examples described herein include systems and methods for parallel upgrades of multiple hosts within an on-premises, multi-rack SDDC. Upgrade instructions can be provided to a POP VM of the SDDC, which can call an API of a pod service executing on a control plane VM in the SDDC. For example, the pod service can generate a set of instructions for performing cluster upgrades in parallel. In one example, the pod service provides the instructions to a lifecycle management service executing on a VM of the SDDC, such as the control plane VM. In some examples, the lifecycle management service can add a spare host to a cluster of hosts in order to begin upgrading one or more hosts in that cluster. This can be performed in parallel, using the spare hosts from different racks simultaneously.

(9) FIG. 1 provides a flowchart of an example method for parallel upgrades of multiple hosts within an on-premises, multi-rack SDDC. At stage **110**, a POP VM executing on the SDDC receives upgrade instructions from a cloud backend. The POP VM can be deployed at a strategic location in the SDDC infrastructure to optimize network performance and reduce latency by providing a local presence for network services. It can act as a gateway or hub for network traffic, such as instructions provided by a cloud backend. The POP VM can provide services such as caching, load balancing, and content delivery to optimize network traffic and improve the user experience. It can also provide security functions such as intrusion detection and prevention, firewalls, and VPN gateways to secure network traffic.

(10) The POP VM can poll a cloud backend for instructions. In some examples, the POP VM polls at regular intervals, but it can also poll in response to a notification received at the POP VM. The cloud backend can be a remote server or group of servers that provides backend services to the SDDC. In some examples, the cloud backend is operated by a third party different from the customer running the SDDC itself.

(11) The instructions received at stage **110** can include instructions to upgrade the SDDC. For example, the instructions can identify specific software and/or firmware upgrades to be applied. In some examples, the instructions include one or more upgrade files. In another example, the instructions provide a location, such as a uniform resource locator (“URL”), where the upgrade files are stored. The POP VM, or any other appropriate component of the SDDC, can then retrieve the upgrade files from the URL. The instructions can also include an identification of specific components in the SDDC to be upgraded. For example, the instructions can identify particular racks, clusters, and hosts to be upgrade. They can also map each component to the specific upgrade to be applied. The instructions can also specify an upgrade order, including instructions to perform certain upgrades in parallel by applying them at the same time.

(12) At stage **120**, the POP VM can call a first API of a pod service executing on a control plane VM of the SDDC. The control plane VM can be a VM that is responsible for managing and controlling the networking and compute resources of the SDDC. It can provide the intelligence and orchestration necessary to manage and automate the deployment, configuration, and operation of virtualized resources. The control plane VM serves as the brain of the SDDC, responsible for managing the virtual network functions, virtual storage, and virtual compute resources. The control plane VM can run specialized software that provides a variety of functions, including resource allocation, load balancing, routing, and security. It communicates with other components in the SDDC, such as the VMs and hypervisors that manage them, to provide a comprehensive view of the virtualized infrastructure and ensure that resources are used efficiently.

(13) The control plane VM can include a pod service, which can be a service executing on the control plane VM that can perform various functionality. In one example, the pod service includes API functionality such that another entity—in this example, the POP VM—can invoke one or more APIs to carry out the instructions received at stage **110**. In one example, the POP VM parses the instructions from stage **110** into multiple instructions to be performed in parallel. In those examples, the POP VM invokes multiple APIs of the pod service at stage **120** to communicate the parallel instructions. In another example, the POP VM invokes a single API to provide the

instructions as one package, and the pod service is responsible for breaking down the instruction set into discrete, parallel instructions.

(14) At stage **130**, the control plane VM can generate a set of instructions for performing cluster upgrades in parallel. The instructions can identify a plurality of clusters of the SDDC, such as multiple clusters operating on separate physical racks of the datacenter. In some examples, this stage is performed by the pod service of stage **120**. In another example, the pod service invokes parallel APIs of a lifecycle management service that can also be executing on the control plane VM. The lifecycle management service can include functionality for performing upgrades, either directly or by instructing other components of the SDDC to perform the upgrades.

(15) Stage **140** can include performing a parallel upgrade of at least two of the identified clusters located on separate racks. This stage can include software or firmware upgrades, or both. In some examples, the lifecycle management service can perform software upgrades directly. This can include moving spare hosts to the clusters to be upgraded. In some on-premises data centers, each physical rack ships with only one spare host. This host can be added to a cluster to be upgraded, allowing one or more other hosts of that cluster to be taken offline while the spare host temporarily takes over any available duties. In an example where the SDDC includes three separate racks, this stage can include adding a spare host to three clusters of the SDDC, with each cluster executing on a different rack relative to the other clusters. This can allow each cluster to be upgraded simultaneously, utilizing the spare hosts as efficiently as possible rather than performing the upgrades serially.

(16) In some examples, the lifecycle management service can instruct other components to perform firmware upgrades. An example of such a component is a hardware security module executing on a manager VM. In some examples, a manager VM is located on each cluster and is configured to apply firmware upgrades. The lifecycle management service can therefore instruct multiple hardware security modules on the parallel clusters to perform firmware upgrades simultaneously. In some examples, these firmware upgrades occur at the same time that the software upgrades occur, applying to the same offline hosts. This process is explained in more detail with respect to FIG. 2.

(17) FIG. 2 provides a sequence diagram of an example method for parallel upgrades of multiple hosts within an on-premises, multi-rack SDDC. At stage **202**, a cloud backend can provide upgrade instructions to the POP VM of the SDDC, as also explained with respect to stage **110** of FIG. 1. The POP VM can poll the cloud backend for instructions. In some examples, the POP VM polls at regular intervals, but it can also poll in response to a notification received at the POP VM. The cloud backend can be a remote server or group of servers that provides backend services to the SDDC. In some examples, the cloud backend is operated by a third party different from the customer running the SDDC itself.

(18) The instructions received at stage **202** can include instructions to upgrade the SDDC. For example, the instructions can identify specific software and/or firmware upgrades to be applied. In some examples, the instructions include one or more upgrade files. In another example, the instructions provide a location, such as a URL, where the upgrade files are stored. The POP VM, or any other appropriate component of the SDDC, can then retrieve the upgrade files from the URL. The instructions can also include an identification of specific components in the SDDC to be upgraded. For example, the instructions can identify particular racks, clusters, and hosts to be upgrade. They can also map each component to the specific upgrade to be applied. The instructions can also specify an upgrade order, including instructions to perform certain upgrades in parallel by applying them at the same time.

(19) At stage **204**, the POP VM can invoke parallel APIs of the control plane VM. In some examples, this stage includes invoking a single API of the control plane VM, such as a pod service executing on the control plane VM. The pod service can be a service executing on the control plane VM that can perform various functionality. In one example, the pod service includes API functionality such that another entity, such as the POP VM, can invoke one or more APIs to carry

out the instructions received at stage **202**. In one example, the POP VM parses the instructions from stage **202** into multiple instructions to be performed in parallel. In those examples, the POP VM invokes multiple APIs of the pod service at stage **204** to communicate the parallel instructions. (20) In another example, the POP VM invokes a single API to provide the instructions as one package, and the pod service is responsible for breaking down the instruction set into discrete, parallel instructions. In that example, the pod service can then generate the parallel API calls to a lifecycle management service of the control plane VM. As explained with respect to stage **130** of FIG. **1**, the control plane VM can generate a set of instructions for performing cluster upgrades in parallel. The instructions can identify a plurality of clusters of the SDDC, such as multiple clusters operating on separate physical racks of the datacenter. In some examples, this stage is performed by the pod service. In another example, the pod service invokes parallel APIs of a lifecycle management service that can also be executing on the control plane VM. The lifecycle management service can include functionality for performing upgrades, either directly or by instructing other components of the SDDC to perform the upgrades. In FIG. **2**, the pod service and lifecycle management service are collectively referred to as the control plane VM.

(21) At stage **206**, the control plane VM can add a spare host to rack 1, which can be one of several physical racks in the datacenter. Typical on-premises, multi-rack datacenters ship with multiple racks that include one spare host on each rack. While certain cloud-based datacenters do not include limitations on the number of spare hosts that can be utilized, an on-premises datacenter must include the spare hosts within the physical rack in order to utilize it. To save on customer costs, each rack typically includes only one spare host. At stage **206**, the control plane can identify a cluster of rack 1 to be upgraded and add the spare host of rack 1 to that cluster. Similarly, at stage **208**, the control plane VM can add the spare host of rack 2 to a cluster to be upgrade in rack 2. Stages **206** and **208** can be performed simultaneously in some examples.

(22) At stage **210**, the control plane VM can cause data to be moved from one or more hosts of the target cluster to the spare host that was added to the cluster at stage **206**. In some examples, the spare host can take over all duties being performed by another host in the target cluster, such that the other host can be taken offline without any degradation of performance or services to the customer. At stage **212**, the control plane VM can cause data to be moved from a different host of a cluster on rack 2 to the spare host added to that cluster at stage **208**. Stages **210** and **212** can be performed simultaneously, allowing multiple spare hosts to be added to multiple clusters on different racks at the same time. Previously this was not possible—instead, only one host from one rack could be utilized at a time, even if the datacenter included multiple racks that each included a spare host. This led to inefficiencies as upgrades were performed serially instead of in parallel.

(23) At stages **214** and **216**, the control plane VM can cause at least one host on each of rack 1 and rack 2 to be taken offline. For example, the hosts that were migrated to the spare hosts on each rack can be the ones taken offline, meaning that they do not interface with the other components of the SDDC and may be powered off or restarted. Because the data from these hosts is migrated to the spare hosts, the spare hosts can provide seamless operation of the SDDC even when the original hosts are taken offline. In some examples, stages **214** and **216** are performed at the same time, also referred to as being performed in parallel.

(24) While the hosts from rack 1 and rack 2 are offline, at stages **218** and **220** the control plane VM upgrades the software of those hosts. In some examples, the lifecycle management service can perform software upgrades directly. For example, the lifecycle management service can send one or more software upgrade packages to each of the offline hosts and instruct them to apply the upgrades. The instructions can also include instructions to reboot the offline hosts. While these instructions are provided separately to the different offline hosts, they can be sent simultaneously (or at least sent to both offline hosts before any of those hosts have completed their upgrade), thereby providing parallel software upgrades.

(25) At stage **222**, the control plane VM can initiate firmware upgrades on the offline hosts. For

example, stage **222** can include providing parallel upgrade instructions to hardware security modules associated with relevant clusters. An example hardware security module can execute on a manager VM associated with each cluster. The lifecycle management service can therefore instruct multiple hardware security modules of the clusters containing the offline hosts to perform firmware upgrades simultaneously. In some examples, these firmware upgrade instructions occur while the software upgrades occur, applying to the same offline hosts. In other examples, the firmware instructions are provided earlier in the method, such as before stage **218**.

(26) Hardware security modules can be specialized hardware devices used for secure key management, cryptographic operations, and secure storage of sensitive information such as passwords, encryption keys, and digital certificates within the SDDC. In some examples, hardware security modules are integrated into an appliance, such as the manager VM or VMWARE VXRAIL, to provide additional security for the SDDC. These hardware security modules are designed to protect the appliance against various security threats, including unauthorized access, data breaches, and malware attacks. They can provide secure key storage and management, which is essential for many security applications such as digital signing, authentication, and encryption. They can also provide secure boot capabilities, which ensures that the manager VM boots up only with authorized software and firmware.

(27) The hardware security modules can perform firmware upgrades of the offline hosts at stages **224** and **226**. In some examples, these stages are performed at the same time or at least in an overlapping manner. In some examples, stages **224** and **226** are performed before the software upgrades at stages **218** and **220**.

(28) After the software and firmware upgrades are complete, the control plane VM can bring the upgraded hosts back online at stages **228** and **230**, respectively. At stages **232** and **234**, the spare hosts can migrate their data back to the original hosts, which have now been upgraded. The spare hosts can then be used to upgrade other hosts in the same clusters or be moved to new clusters to assist with upgrades of hosts in those clusters.

(29) FIG. 3 provides a flowchart of an example method for parallel upgrades of multiple hosts within an on-premises, multi-rack SDDC. The example method begins at stage **302**, where a data plane upgrade is invoked. This stage can include the POP VM receiving upgrade instructions from a cloud backend and invoking parallel APIs at a control plane VM, for example.

(30) At stage **304**, the control plane VM can determine whether all clusters are upgraded. If so, the process can end at stage **328**. If not, the method can continue to stage **306** and perform SDDC level initialization as necessary. SDDC initialization can include initializing or setting up any aspect of the SDDC to perform upgrades. In one example, this stage includes powering on one or more spare hosts in the SDDCs. For example, in a three-rack SDDC where each rack includes one spare host, the spare hosts can remain dormant in a powered-down state until they are needed. This saves on energy costs. When they are needed, the spare hosts can be powered up and initialized as necessary, such as by installing software or firmware, migrating data, and applying settings or policies to the spare hosts.

(31) If stage **306** is determined to be unsuccessful at stage **308**, then the method can proceed to stage **310** where the failure is handled and the system awaits instructions for further actions. If stage **308** determines that the initialization was successful, then the method can proceed to stage **312**. Stages **306** and **308** can be performed by the control plane VM, for example.

(32) At stage **312**, SDDC level pre-hooks are orchestrated by the control plane VM. These pre-hooks can include, for example, setting up a depot for the lifecycle management service; deleting data in an old depot, if one exists; clearing out the cache; and applying any lifecycle management service global policies. The pre-hooks can be executed only once at the SDDC level, such that the method does not require repeat execution when multiple clusters are being upgraded.

(33) If stage **312** is determined to be unsuccessful at stage **314**, then the method can proceed to stage **310** where the failure is handled, and the system awaits instructions for further actions. If

stage **314** determines that the initialization was successful, then the method can proceed to stage **316**.

(34) At stage **316**, the control plane VM can upgrade a first set of clusters in the SDDC. This can include, for example, the lifecycle management service can perform software upgrades directly. For example, the lifecycle management service can send one or more software upgrade packages to each of the offline hosts and instruct them to apply the upgrades. The instructions can also include instructions to reboot the offline hosts. While these instructions are provided separately to the different offline hosts, they can be sent simultaneously (or at least sent to both offline hosts before any of those hosts have completed their upgrade), thereby providing parallel software upgrades.

(35) Stage **316** can also include firmware upgrades. For example, stage **316** can include providing parallel upgrade instructions to hardware security modules associated with relevant clusters. An example hardware security module can execute on a manager VM associated with each cluster. The lifecycle management service can therefore instruct multiple hardware security modules of the clusters containing the offline hosts to perform firmware upgrades simultaneously. In some examples, these firmware upgrade instructions occur while the software upgrades occur, applying to the same offline hosts.

(36) Stage **316** can also include bringing the offline hosts back online and migrating data from the spare hosts back to the original, now-upgraded hosts. The spare hosts can also be used to migrate any other hosts of the clusters, such that stage **316** includes upgrading all hosts within the first set of clusters.

(37) If stage **316** is determined to be unsuccessful at stage **318**, then the method can proceed to stage **310** where the failure is handled, and the system awaits instructions for further actions. If stage **318** determines that the initialization was successful, then the method can proceed to stage **320**.

(38) Stage **320** can include repeating the process of step **316** for a second, third, or any other number of cluster sets. In an example where an SDDC includes three racks, each having five clusters, then each “set” of clusters can include one cluster from each rack. This will result in five sets of clusters that are upgraded in stages **316-320**.

(39) If stage **320** is determined to be unsuccessful at stage **322**, then the method can proceed to stage **310** where the failure is handled, and the system awaits instructions for further actions. If stage **322** determines that the initialization was successful, then the method can proceed to stage **324**.

(40) At stage **324**, the SDDC can perform post-hooks, which can include various checks and adjustments after upgrading is complete. This can include confirming that each host is upgraded and that the inventories for the hosts have been synchronized, performing a disk format conversion change, and performing any other required configuration changes. If stage **324** is determined to be unsuccessful at stage **326**, then the method can proceed to stage **310** where the failure is handled, and the system awaits instructions for further actions. If stage **326** determines that the initialization was successful, then the method can end at stage **328**. Stage **326** can be performed by the control plane VM, for example.

(41) FIG. 4 provides a flowchart of an example method for parallel upgrades of multiple hosts within an on-premises, multi-rack SDDC. At stage **402**, the SDDC can initialize clusters 1-n within the SDDC. In this example, clusters 1-n are distributed amongst at least three separate racks. This stage can include hardware setup, such as configuring the servers, switches, and storage devices of the SDDC. It can also include network configuration, such as setting up network adapters, VLANs, and other network settings required for communication between the hosts and VMs. It can further include cluster creation, such as installing hypervisor software on each host and configuring them to work together in a cluster. This stage can also include configuring storage infrastructure and deploying VMs to the clusters.

(42) At stage **404**, the spare hosts from each rack can be added to various clusters. In the examples

described previously, each rack included one spare host. In those examples, the one spare host from each rack can be added to one cluster located on the respective rack. However, the systems and methods described herein are scalable such that any number of spare hosts can be utilized at the same time. For example, in a three-rack SDDC where each rack includes two spare hosts, both spare hosts of each rack can be added to clusters within that rack. The methods described above, such as the methods of FIGS. 1-3, can utilize multiple spare hosts to further increase the parallelism of the upgrade process. For example, stage 218 of FIG. 2 describes upgrading software on an offline host, where the offline host has been migrated to a spare host. In an example where the rack for that host includes two spare hosts, stage 218 can include upgrading software on two offline hosts in that rack, where the offline hosts have been migrated to the two spare hosts.

(43) At stage 406, the SDDC can run pre-hooks on all clusters to be upgraded. These pre-hooks can be performed in parallel on the various clusters. Pre-hooks can include scripts or commands that are executed before a cluster operation is performed. These pre-hooks are designed to automate tasks that need to be performed before a cluster operation, such as adding or removing a host, updating the cluster configuration, or performing maintenance tasks. Some examples of pre-hooks include checking the health status of a host and configuring settings of the host, such as network settings, firewall rules, and storage settings.

(44) At stage 408, the SDDC can invoke a pod service API to perform data plane pre-checks. These pre-checks can include tests to verify that the data plane of the SDDC is ready for use. The data plane is the part of the SDDC that handles the actual data traffic, such as VM network traffic or storage I/O. The purpose of data plane pre-checks is to ensure that the data plane is properly configured and functioning correctly before VMs or other workloads are deployed onto the SDDC. Invocation of the pod service using APIs is described in more detail with respect to stage 204 of FIG. 2, for example.

(45) At stage 410, the SDDC can perform data plane upgrades on the clusters in parallel. These upgrades can include software and firmware upgrades in some examples. For example, the control plane VM upgrades the software of those hosts. In some examples, the lifecycle management service can perform software upgrades directly. For example, the lifecycle management service can send one or more software upgrade packages to each of the offline hosts and instruct them to apply the upgrades. The instructions can also include instructions to reboot the offline hosts. While these instructions are provided separately to the different offline hosts, they can be sent simultaneously (or at least sent to both offline hosts before any of those hosts have completed their upgrade), thereby providing parallel software upgrades.

(46) Similarly, the control plane VM can initiate firmware upgrades on the offline hosts. For example, stage 410 can include providing parallel upgrade instructions to hardware security modules associated with relevant clusters. An example hardware security module can execute on a manager VM associated with each cluster. The lifecycle management service can therefore instruct multiple hardware security modules of the clusters containing the offline hosts to perform firmware upgrades simultaneously. In some examples, these firmware upgrade instructions occur while the software upgrades occur, applying to the same offline hosts. In other examples, the firmware instructions are provided earlier in the method.

(47) At stage 412, the pod service can be invoked through an API call to perform data plane post-checks. The post-checks are tests that are performed after a workload has been deployed onto the SDDC to ensure that the data plane is functioning correctly. The purpose of data-plane post-checks is to verify that the data plane is performing as expected after a workload has been deployed. This helps ensure that the workload is operating correctly and that there are no issues with the SDDC infrastructure. These post-checks typically include tests of network connectivity, storage connectivity, security settings, and VM performance.

(48) At stage 414, the SDDC can run post-hooks on the clusters in parallel. Post-hooks can be scripts or commands that are executed after a cluster operation has been performed. These hooks

are designed to automate tasks that need to be performed after a cluster operation, such as adding or removing a host, updating the cluster configuration, or performing maintenance tasks. These post-hooks can include, for example, validating the configuration of the cluster, triggering additional tasks such as updating a backup schedule or performing additional maintenance tasks, checking the health status of the cluster to ensure it is functioning appropriately, and sending notifications.

(49) At stage **416**, the method can include removing the spare hosts from the clusters that were upgraded. This stage can be performed in parallel, such that the spare hosts are removed at approximately the same time. In an example system having three racks with two spare hosts per rack, the upgrade process can utilize all spare hosts as explained above. At stage **416**, all sick spare hosts can be removed at the same time, such as by migrating their data back to the original hosts. The process can then end at stage **418** or repeat as needed until all relevant clusters and hosts are upgraded.

(50) FIG. 5 provides an illustration of an example system for performing the methods described herein. FIG. 5 depicts an SDDC **520** that includes three physical racks **522**, **524**, **526**. Each rack **522**, **524**, **526** includes various clusters that each include one or more hosts. Each rack **522**, **524**, **526** also includes a spare host, although in some examples each rack can include more than one spare host. The SDDC **520** also includes a POP VM **530**, a control plane VM **540**, and multiple manager VMs **550**, **554**.

(51) The POP VM **530** can be deployed at a strategic location in the SDDC infrastructure to optimize network performance and reduce latency by providing a local presence for network services. It can act as a gateway or hub for network traffic, such as instructions provided by a cloud backend. The POP VM **530** can provide services such as caching, load balancing, and content delivery to optimize network traffic and improve the user experience. It can also provide security functions such as intrusion detection and prevention, firewalls, and VPN gateways to secure network traffic.

(52) The POP VM **530** can include a task framework service **531** and an LCM service **532**, which can each be software services executing on the POP VM **530**. The task framework service **531** can be a set of rules and procedures for managing and executing tasks within the POP VM **530**. In some examples, the task framework service **531** can poll a cloud backend **512** for tasks. For example, a data plane upgrade job can be a task, or set of tasks, provided by the cloud backend **512** to the task framework service **531** of the POP VM **530**. In some examples, the POP VM **530** polls at regular intervals, but it can also poll in response to a notification received at the POP VM **530**. The cloud backend **512** can be a remote server or group of servers that provides backend services to the SDDC. In some examples, the cloud backend **512** is operated by a third party different from the customer running the SDDC itself.

(53) The upgrade job received from the cloud backend **512** can be created at a fleet management service **510**. In some examples, the fleet management service **510** is managed by the same entity that operates the cloud backend **512**. For example, a third-party that provides SDDC infrastructure and management can operate the fleet management service **510**. The fleet management service **510** can generate the data plane upgrade job and transmit it to one or more cloud backends **512** that support various SDDCs, including SDDC **520**.

(54) Turning back to the POP VM **530**, the task framework service **531** can receive the data plane upgrade job and schedule it using a lifecycle management service **532**. The lifecycle management service **532** of the POP VM **530** can determine when and how to apply the upgrades using information about the current state of the SDDC **520**. For example, it can determine that the upgrades have not yet been applied to certain clusters or hosts within the SDDC **520**. It can also determine how many spare hosts are available and whether parallel upgrades can be performed. It can invoke a pod service **541** of a control plane VM **540** using API calls or other communication techniques, instructing the pod service **541** to take actions to carry out the upgrade plane.

(55) The pod service **541** can be a service executing on the control plane VM **540** that can perform

various functionality. In one example, the pod service **541** includes API functionality such that another entity, such as the POP VM **530**, can invoke one or more APIs to carry out the instructions received from the cloud backend **512**. In one example, the POP VM **530** parses the instructions into multiple instructions to be performed in parallel. In those examples, the POP VM **530** invokes multiple APIs of the pod service **541** to communicate the parallel instructions.

(56) In another example, the POP VM **530** invokes a single API to provide the instructions as one package, and the pod service **541** is responsible for breaking down the instruction set into discrete, parallel instructions. In that example, the pod service **541** can then generate the parallel API calls to a lifecycle management service **542** of the control plane VM **540**. The control plane VM **540** can generate a set of instructions for performing cluster upgrades in parallel. The instructions can identify a plurality of clusters of the SDDC **520**, such as multiple clusters operating on separate physical racks of the datacenter. In some examples, this stage is performed by the pod service **541**. In another example, the pod service **541** invokes parallel APIs of a lifecycle management service **542** executing on the control plane VM **540**. The lifecycle management service **542** can include functionality for performing upgrades, either directly or by instructing other components of the SDDC **520** to perform the upgrades.

(57) The lifecycle management service **542** of the control plane VM **542** can perform parallel software upgrades of the clusters in the manner explained in various examples above. In some examples, the lifecycle management service can perform software upgrades directly. For example, the lifecycle management service can send one or more software upgrade packages to each of the offline hosts and instruct them to apply the upgrades. The instructions can also include instructions to reboot the offline hosts. While these instructions are provided separately to the different offline hosts, they can be sent simultaneously (or at least sent to both offline hosts before any of those hosts have completed their upgrade), thereby providing parallel software upgrades.

(58) The lifecycle management service **542** of the control plane VM **542** can also initiate firmware upgrades on the offline hosts. For example, it can provide parallel upgrade instructions to hardware security modules **552**, **554** associated with relevant clusters. An example hardware security module **552**, **554** can execute on a manager VM **550**, **554** associated with each cluster. The lifecycle management service **542** can therefore instruct multiple hardware security modules **552**, **554** of the clusters containing the offline hosts to perform firmware upgrades simultaneously. In some examples, these firmware upgrade instructions occur while the software upgrades occur, applying to the same offline hosts.

(59) Hardware security modules **552**, **556** can be specialized hardware devices used for secure key management, cryptographic operations, and secure storage of sensitive information such as passwords, encryption keys, and digital certificates within the SDDC **520**. In some examples, hardware security modules **552**, **556** are integrated into an appliance, such as the manager VM or VMWARE VXRAIL, to provide additional security for the SDDC **520**. These hardware security modules **552**, **556** are designed to protect the appliance against various security threats, including unauthorized access, data breaches, and malware attacks. They can provide secure key storage and management, which is essential for many security applications such as digital signing, authentication, and encryption. They can also provide secure boot capabilities, which ensures that the manager VM **550**, **556** boots up only with authorized software and firmware.

(60) FIG. **6** provides an illustration of a simplified SDDC that can perform the various methods described herein. Specifically, FIG. **6** illustrates a virtual data center provided as an abstraction of underlying physical-data-center hardware components. In some examples, the virtualized data center is considered an SDDC. The virtualized data center can include various components described herein, such as the POP VM, control plane VM, manager VMs, and server racks. A physical data center **602** is shown below a virtual-interface plane **604**. The physical data center consists of a virtual-infrastructure management server (“VI-management-server”) **606** and any of various different computers, such as PCs **608**, on which a virtual-data-center management interface

can be displayed to system administrators and other users. The interface can be a software-based control system, such as VMWARE VCLOUD DIRECTOR.

(61) The physical data center additionally includes a number of server computers, such as server computer **610**, that are coupled together by local area networks, such as local area network **612** that directly interconnects server computer **610** and **614-620** and a mass-storage array **622**. The physical data center shown in FIG. **6** includes three local area networks **612**, **624**, and **626** that each directly interconnects a bank of eight servers and a mass-storage array. The individual server computers, such as server computer **610**, each includes a virtualization layer and runs multiple VMs. Different physical data centers can include many different types of computers, networks, data-storage systems and devices connected according to many different types of connection topologies. The virtual-data-center abstraction layer **604**, a logical abstraction layer shown by a plane in FIG. **6**, abstracts the physical data center to a virtual data center comprising one or more resource pools, such as resource pools **630-632**, one or more virtual data stores, such as virtual data stores **634-636**, and one or more virtual networks. In certain implementations, the resource pools abstract banks of physical servers directly interconnected by a local area network.

(62) The virtual-data-center management interface allows provisioning and launching of VMs with respect to resource pools, virtual data stores, and virtual networks, so that virtual-data-center administrators need not be concerned with the identities of physical-data-center components used to execute particular VMs. Furthermore, the VI-management-server includes functionality to migrate running VMs from one physical server to another in order to optimally or near optimally manage resource allocation, provide fault tolerance, and high availability. It can migrate VMs to most effectively utilize underlying physical hardware resources, to replace VMs disabled by physical hardware problems and failures, and to ensure that multiple VMs supporting a high-availability virtual appliance are executing on multiple physical computer systems. This ensures the services provided by the virtual appliance are continuously accessible, even when one of the multiple virtual appliances becomes compute bound, data-access bound, suspends execution, or fails. Thus, the virtual data center layer of abstraction provides a virtual-data-center abstraction of physical data centers to simplify provisioning, launching, and maintenance of VMs and virtual appliances as well as to provide high-level, distributed functionalities that involve pooling the resources of individual physical servers and migrating VMs among physical servers to achieve load balancing, fault tolerance, and high availability.

(63) The virtual data center provided by a service provider can be configured through a control system interface displayed on a user device, such as a computer or smartphone. In some examples, both the service provider and the tenants of that service provider can access certain configuration information through a GUI associated with the control system.

(64) The SDDC of FIG. **6** can include multiple SDDC instances executing within the SDDC. For example, the SDDC can include an SDDC manager instance that operates on a standalone SDDC stack within the overall SDDC. Similarly, the SDDC can include an SDDC workflow instance that operates on another SDDC stack. The SDDC can include multiple SDDC manager instances and SDDC workflow instances, as described above.

(65) Although examples described herein include systems and methods for dynamic VM provisioning across different cloud service providers, the same systems and methods can be applied for dynamic VM provisioning within a single cloud service provider. The examples or statements above relating to provisioning across cloud service providers are not intended to be limited regarding whether such provisioning is across different cloud service providers or within a single cloud service provider.

(66) Other examples of the disclosure will be apparent to those skilled in the art from consideration of the specification and practice of the examples disclosed herein. Though some of the described methods have been presented as a series of steps, it should be appreciated that one or more steps can occur simultaneously, in an overlapping fashion, or in a different order. The order of steps

presented are only illustrative of the possibilities and those steps can be executed or performed in any suitable fashion. Moreover, the various features of the examples described here are not mutually exclusive. Rather any feature of any example described here can be incorporated into any other suitable example. It is intended that the specification and examples be considered as exemplary only, with a true scope and spirit of the disclosure being indicated by the following claims.

Claims

1. A method for upgrading a multi-rack, on-premises software defined data center (SDDC), comprising: receiving upgrade instructions at a point-of-presence (POP) virtual machine (VM) executing on the SDDC; generating, by a pod service, a set of instructions for performing cluster upgrades in parallel, wherein the set of instructions identifies a plurality of clusters of the SDDC located on separate racks; and performing a parallel upgrade of at least two of the identified plurality of clusters located on separate racks, wherein performing the parallel upgrade comprises: adding a spare host to each of the identified plurality of clusters, wherein each rack of the SDDC includes at least one spare host; transferring data from at least one host on each of the identified plurality of clusters to the spare host added to that cluster; taking the at least one host on each of the identified plurality of clusters offline; upgrading at least one of software and firmware on the offline hosts; bringing the offline hosts online; and removing the spare hosts from the identified plurality of clusters.
2. The method of claim 1, further comprising calling, by the POP VM, a first application programming interface (API) of the pod service executing on a control plane VM of the SDDC, the call including instructions to upgrade the plurality of clusters of the SDDC in parallel.
3. The method of claim 1, wherein upgrading comprises upgrading both software and firmware, and wherein the upgrading is performed in parallel on the offline hosts.
4. The method of claim 1, wherein upgrading firmware on the offline hosts is performed by hardware security modules executing on each of the identified plurality of clusters, respectively.
5. The method of claim 4, wherein the hardware security modules upgrade respective firmware based on instructions received from a lifecycle management service executing on a control plane VM of the SDDC.
6. The method of claim 1, wherein upgrading software on the offline hosts is performed by a lifecycle management service executing on a control plane VM of the SDDC.
7. The method of claim 1, wherein each rack includes only one spare host.
8. A non-transitory, computer-readable medium containing instructions that, when executed by a hardware-based processor, performs stages for upgrading a multi-rack, on-premises software defined data center (SDDC), the stages comprising: receiving upgrade instructions at a point-of-presence (POP) virtual machine (VM) executing on the SDDC; generating, by a pod service, a set of instructions for performing cluster upgrades in parallel, wherein the set of instructions identifies a plurality of clusters of the SDDC located on separate racks; and performing a parallel upgrade of at least two of the identified plurality of clusters located on separate racks, wherein performing the parallel upgrade comprises: adding a spare host to each of the identified plurality of clusters, wherein each rack of the SDDC includes at least one spare host; transferring data from at least one host on each of the identified plurality of clusters to the spare host added to that cluster; taking the at least one host on each of the identified plurality of clusters offline; upgrading at least one of software and firmware on the offline hosts; bringing the offline hosts online; and removing the spare hosts from the identified plurality of clusters.
9. The non-transitory, computer-readable medium of claim 8, the stages further comprising calling, by the POP VM, a first application programming interface (API) of the pod service executing on a control plane VM of the SDDC, the call including instructions to upgrade the plurality of clusters

of the SDDC in parallel.

10. The non-transitory, computer-readable medium of claim 8, wherein upgrading comprises upgrading both software and firmware, and wherein the upgrading is performed in parallel on the offline hosts.

11. The non-transitory, computer-readable medium of claim 8, wherein upgrading firmware on the offline hosts is performed by hardware security modules executing on each of the identified plurality of clusters, respectively.

12. The non-transitory, computer-readable medium of claim 11, wherein the hardware security modules upgrade respective firmware based on instructions received from a lifecycle management service executing on a control plane VM of the SDDC.

13. The non-transitory, computer-readable medium of claim 8, wherein upgrading software on the offline hosts is performed by a lifecycle management service executing on a control plane VM of the SDDC.

14. The non-transitory, computer-readable medium of claim 8, wherein each rack includes only one spare host.

15. A system for upgrading a multi-rack, on-premises software defined data center (SDDC), comprising: a memory storage including a non-transitory, computer-readable medium comprising instructions; and a computing device including a hardware-based processor that executes the instructions to carry out stages comprising: receiving upgrade instructions at a point-of-presence (POP) virtual machine (VM) executing on the SDDC; generating, by a pod service, a set of instructions for performing cluster upgrades in parallel, wherein the set of instructions identifies a plurality of clusters of the SDDC located on separate racks; and performing a parallel upgrade of at least two of the identified plurality of clusters located on separate racks, wherein performing the parallel upgrade comprises: adding a spare host to each of the identified plurality of clusters, wherein each rack of the SDDC includes at least one spare host; transferring data from at least one host on each of the identified plurality of clusters to the spare host added to that cluster; taking the at least one host on each of the identified plurality of clusters offline; upgrading at least one of software and firmware on the offline hosts; bringing the offline hosts online; and removing the spare hosts from the identified plurality of clusters.

16. The system of claim 15, the stages further comprising calling, by the POP VM, a first application programming interface (API) of the pod service executing on a control plane VM of the SDDC, the call including instructions to upgrade the plurality of clusters of the SDDC in parallel.

17. The system of claim 15, wherein upgrading comprises upgrading both software and firmware, and wherein the upgrading is performed in parallel on the offline hosts.

18. The system of claim 15, wherein upgrading firmware on the offline hosts is performed by hardware security modules executing on each of the identified plurality of clusters, respectively.

19. The system of claim 18, wherein the hardware security modules upgrade respective firmware based on instructions received from a lifecycle management service executing on a control plane VM of the SDDC.

20. The system of claim 15, wherein upgrading software on the offline hosts is performed by a lifecycle management service executing on a control plane VM of the SDDC.
