

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12399706
Kind Code	B2
Date of Patent	August 26, 2025
Inventor(s)	Kirmse; Daniel

Cloud version management for legacy on-premise application

Abstract

Various examples are directed to systems and methods for upgrading a cloud-implemented application. A cloud service may receive a request from a user group to access the application. The cloud service may access consumer context data comprising a plurality of context properties of the user group and may access a rollout strategy map comprising a first rollout record. The cloud service may compare the plurality of context properties of the first user group to first selector data indicated by a first rollout record. Based on the comparing, the cloud service may add the first version of the application to a list of permissible versions for the first user group.

Inventors:	Kirmse; Daniel (Berlin, DE)
Applicant:	SAP SE (Walldorf, DE)
Family ID:	1000008781664
Assignee:	SAP SE (Walldorf, DE)
Appl. No.:	18/398801
Filed:	December 28, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240126540 A1	Apr. 18, 2024

Related U.S. Application Data

continuation parent-doc US 17484667 20210924 US 11922163 child-doc US 18398801

Publication Classification

Int. Cl.: G06F8/65 (20180101); **G06F8/656** (20180101); **G06F8/71** (20180101); **G06F9/48** (20060101); **G06F9/54** (20060101); **H04L67/10** (20220101)

U.S. Cl.:

CPC **G06F8/656** (20180201); **G06F9/48** (20130101); **H04L67/10** (20130101); G06F8/65 (20130101); G06F8/71 (20130101)

Field of Classification Search

CPC: G06F (8/656); G06F (8/65); G06F (8/71); G06F (9/48); H04L (67/10)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
9817653	12/2016	Sreedharan et al.	N/A	N/A
10320625	12/2018	Cherumbath	N/A	N/A
10338964	12/2018	Kancharla et al.	N/A	N/A
10387385	12/2018	Mart	N/A	N/A
10496392	12/2018	Shuvali et al.	N/A	N/A
10732951	12/2019	Jayanthi et al.	N/A	N/A
10853146	12/2019	Talayco	N/A	N/A
11113090	12/2020	Wilkinson et al.	N/A	N/A
11372820	12/2021	Harjono et al.	N/A	N/A
11704408	12/2022	Ciubotariu et al.	N/A	N/A
11841731	12/2022	Kirmse	N/A	N/A
12013866	12/2023	Kirmse	N/A	N/A
12026496	12/2023	Kirmse	N/A	N/A
2006/0167947	12/2005	Dunkle	N/A	N/A
2007/0067361	12/2006	Bailey et al.	N/A	N/A
2009/0177646	12/2008	Pham et al.	N/A	N/A
2011/0083138	12/2010	Sivasubramanian et al.	N/A	N/A
2013/0190080	12/2012	Bibbey et al.	N/A	N/A
2013/0263088	12/2012	Hoff et al.	N/A	N/A
2016/0188415	12/2015	Karinta et al.	N/A	N/A
2016/0275019	12/2015	Nam et al.	N/A	N/A
2016/0294614	12/2015	Searle et al.	N/A	N/A
2017/0046181	12/2016	Williams et al.	N/A	N/A
2017/0115978	12/2016	Modi et al.	N/A	N/A
2017/0161053	12/2016	Hayakawa et al.	N/A	N/A
2017/0192873	12/2016	Ozdemir et al.	N/A	N/A
2017/0300311	12/2016	Vasquez Lopez et al.	N/A	N/A
2017/0330195	12/2016	Lange et al.	N/A	N/A
2018/0004792	12/2017	Desai et al.	N/A	N/A
2018/0011699	12/2017	Manthiramoorthy et al.	N/A	N/A
2018/0267796	12/2017	Kennedy et al.	N/A	N/A

2019/0102280	12/2018	Caldato et al.	N/A	N/A
2019/0146772	12/2018	Griffin et al.	N/A	N/A
2019/0317887	12/2018	Wiener et al.	N/A	N/A
2019/0327294	12/2018	Subramani Nadar	N/A	H04L 41/0843
2019/0370471	12/2018	Petratos et al.	N/A	N/A
2020/0019414	12/2019	Byard et al.	N/A	N/A
2020/0145337	12/2019	Keating et al.	N/A	N/A
2020/0228602	12/2019	Spoczynski et al.	N/A	N/A
2020/0241864	12/2019	Duvur	N/A	G06F 8/65
2020/0322226	12/2019	Mishra	N/A	H04L 43/0876
2021/0049267	12/2020	Wu et al.	N/A	N/A
2022/0050674	12/2021	Liljeback et al.	N/A	N/A
2022/0107802	12/2021	Rao et al.	N/A	N/A
2022/0292106	12/2021	Umay et al.	N/A	N/A
2022/0294788	12/2021	Pattar et al.	N/A	N/A
2022/0350642	12/2021	Poddar et al.	N/A	N/A
2022/0365801	12/2021	Kasso et al.	N/A	N/A
2023/0070209	12/2022	Srinivasan et al.	N/A	N/A
2023/0096143	12/2022	Kirmse	N/A	N/A
2023/0098023	12/2022	Kirmse	N/A	N/A
2023/0102514	12/2022	Kirmse	N/A	N/A
2023/0102769	12/2022	Kirmse	N/A	N/A
2023/0103087	12/2022	Kirmse	N/A	N/A
2023/0103223	12/2022	Kirmse	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
103268319	12/2012	CN	N/A
110347440	12/2018	CN	N/A
108595221	12/2020	CN	N/A
3451341	12/2018	EP	N/A
3557408	12/2018	EP	N/A
WO-2013148144	12/2012	WO	N/A
WO-2016028973	12/2015	WO	N/A

OTHER PUBLICATIONS

Katarina Grolinger et al., Data management in cloud environments: NoSQL and NewSQL data stores, 2013, [Retrieved on Mar. 3, 2025]. Retrieved from the internet: <URL: <https://link.springer.com/content/pdf/10.1186/2192-113X-2-22.pdf>> 24 Pages (1-24) (Year: 2013). cited by examiner

“U.S. Appl. No. 17/484,637, Final Office Action mailed Mar. 23, 2023”, 9 pgs. cited by applicant

“U.S. Appl. No. 17/484,637, Non Final Office Action mailed Aug. 2, 2023”, 8 pgs. cited by applicant

“U.S. Appl. No. 17/484,637, Non Final Office Action mailed Nov. 10, 2022”, 9 pgs. cited by applicant

“U.S. Appl. No. 17/484,637, Notice of Allowance mailed May 26, 2023”, 8 pgs. cited by applicant

“U.S. Appl. No. 17/484,637, Preliminary Amendment filed Jun. 20, 2023”, 3 pgs. cited by

applicant

“U.S. Appl. No. 17/484,637, Response filed Feb. 10, 2023 to Non Final Office Action mailed Nov. 10, 2022”, 12 pgs. cited by applicant

“U.S. Appl. No. 17/484,637, Response filed May 4, 2023 to Final Office Action mailed Mar. 23, 2023”, 12 pgs. cited by applicant

“U.S. Appl. No. 17/484,637, Response filed Nov. 2, 2023 to Non Final Office Action mailed Aug. 2, 2023”, 9 pgs. cited by applicant

“U.S. Appl. No. 17/484,667, Notice of Allowance mailed Apr. 12, 2023”, 22 pgs. cited by applicant

“U.S. Appl. No. 17/484,667, Notice of Allowance mailed Jul. 10, 2023”, 8 pgs. cited by applicant

“U.S. Appl. No. 17/484,667, Notice of Allowance mailed Nov. 7, 2023”, 8 pgs. cited by applicant

“U.S. Appl. No. 17/484,667, Preliminary Amendment filed Jun. 20, 2023”, 3 pgs. cited by applicant

“U.S. Appl. No. 17/484,688, Non Final Office Action mailed Nov. 25, 2022”, 20 pgs. cited by applicant

“U.S. Appl. No. 17/484,688, Notice of Allowance mailed May 3, 2023”, 15 pgs. cited by applicant

“U.S. Appl. No. 17/484,688, Notice of Allowance mailed Jul. 31, 2023”, 18 pgs. cited by applicant

“U.S. Appl. No. 17/484,688, Response filed Feb. 27, 2023 to Non Final Office Action mailed Nov. 25, 2022”, 10 pgs. cited by applicant

Afgan, Enis, et al., “CloudLaunch: Discover and Deploy Cloud Applications”, arXiv:1805.04005, [Online]. Retrieved from the Internet: <URL:

<https://arxiv.org/ftp/arxiv/papers/1805/1805.04005.pdf>>, (2018), 9 pgs. cited by applicant

Lyublena, Antova, “Rapid Adoption of Cloud Data Warehouse Technology using Datometry Hyper-Q (Technical White Paper)”, Datometry Inc., [Online]. Retrieved from the Internet:

<<https://dl.acm.org/doi/pdf/10.1145/3183713.3190652>>, (2018), 1-12. cited by applicant

Sharma, Prateek, “Containers and Virtual Machines at Scale: A Comparative Study”, Proceedings of the 17th International Middleware Conference, Article No. 1, [Online]. Retrieved from the

Internet: <<https://dl.acm.org/doi/pdf/10.1145/2988336.2988337>>, (2016), 1-13. cited by applicant

Zurkowski, Bartosz, “Towards Self-organizing Cloud Polyglot Database Systems”, IEEE 13th International Conference on Self-Adaptive and Self-Organizing Systems, [Online] Retrieved from

the internet: <<https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=8780526>>, (2019), 82-87. cited by applicant

“European Application Serial No. 22196329.1, Extended European Search Report mailed Mar. 1, 2023”, 14 pgs. cited by applicant

“U.S. Appl. No. 17/484,704, Non Final Office Action mailed Mar. 31, 2023”, 12 pgs. cited by applicant

“U.S. Appl. No. 17/484,754, Non Final Office Action mailed Apr. 13, 2023”, 30 pgs. cited by applicant

“U.S. Appl. No. 17/484,704, Response filed Jun. 21, 2023 to Non Final Office Action mailed Mar. 31, 2023”, 10 pgs. cited by applicant

“U.S. Appl. No. 17/484,754, Response filed Jul. 6, 2023 to Non Final Office Action mailed Apr. 13, 2023”, 14 pgs. cited by applicant

“U.S. Appl. No. 17/484,704, Final Office Action mailed Sep. 6, 2023”, 14 pgs. cited by applicant

“U.S. Appl. No. 17/484,754, Final Office Action mailed Sep. 21, 2023”, 34 pgs. cited by applicant

“U.S. Appl. No. 17/484,732, Non Final Office Action mailed Oct. 26, 2023”, 27 pgs. cited by applicant

“U.S. Appl. No. 17/484,704, Response filed Dec. 6, 2023 to Final Office Action mailed Sep. 6, 2023”, 11 pgs. cited by applicant

“U.S. Appl. No. 17/484,704, Non Final Office Action mailed Dec. 14, 2023”, 14 pgs. cited by applicant

“U.S. Appl. No. 17/484,754, Response filed Dec. 21, 2023 to Final Office Action mailed Sep. 21,

2023", 12 pgs. cited by applicant
"U.S. Appl. No. 17/484,732, Response filed Jan. 17, 2024 to Non Final Office Action mailed Oct. 26, 2023", 11 pgs. cited by applicant
"U.S. Appl. No. 17/484,667, Corrected Notice of Allowability mailed Jan. 29, 2024", 13 pgs. cited by applicant
"U.S. Appl. No. 17/484,637, Notice of Allowance mailed Feb. 15, 2024", 8 pgs. cited by applicant
"U.S. Appl. No. 17/484,732, Notice of Allowance mailed Feb. 22, 2024", 10 pgs. cited by applicant
"U.S. Appl. No. 17/484,704, Examiner Interview Summary mailed Feb. 28, 2024", 2 pgs. cited by applicant
"U.S. Appl. No. 17/484,704, Response filed Feb. 29, 2024 to Non Final Office Action mailed Dec. 14, 2023", 11 pgs. cited by applicant
"U.S. Appl. No. 17/484,754, Non Final Office Action mailed Mar. 7, 2024", 35 pgs. cited by applicant
"U.S. Appl. No. 17/484,754, Examiner Interview Summary mailed May 23, 2024", 3 pgs. cited by applicant
"U.S. Appl. No. 17/484,754, Response filed Jun. 7, 2024 to Non Final Office Action mailed Mar. 7, 2024", 14 pgs. cited by applicant
"U.S. Appl. No. 17/484,704, Notice of Allowance mailed Jul. 18, 2024", 7 pgs. cited by applicant
"U.S. Appl. No. 17/484,754, Notice of Allowance mailed Jul. 31, 2024", 19 pgs. cited by applicant
"U.S. Appl. No. 17/484,754, Corrected Notice of Allowability mailed Aug. 12, 2024", 2 pgs. cited by applicant
"U.S. Appl. No. 17/484,754, PTO Response to Rule 312 Communication mailed Oct. 10, 2024", 2 pgs. cited by applicant
"U.S. Appl. No. 17/484,754,312 Amendment filed Oct. 8, 2024", 6 pgs. cited by applicant
"U.S. Appl. No. 17/484,704, Notice of Allowance mailed Nov. 22, 2024", 7 pgs. cited by applicant
Pahl, Claus, "Containerization and the PaaS Cloud", IEEE Cloud Computing, vol. 1, No. 3, (May 1, 2015), 24-31. cited by applicant
Shah, Jay, "Building Modern Clouds: Using Docker, Kubernetes and Google Cloud Platform", IEEE 9th Annual Computing and Communication Workshop and Conference, (Jan. 7, 2019), 0184-0189. cited by applicant
"European Application Serial No. 22196329.1 Communication pursuant to Rule 94(3) EPC filed on Jun. 6, 2025", 11 pgs. cited by applicant
"Deployment environment", [Online]. Retrieved from the Internet: URL: https://en.wikipedia.org/w/index.php?title=Deployment_environment&oldid=1023026859, (Accessed on May 13, 2021), 5 pgs. cited by applicant

Primary Examiner: Rivera; Anibal

Attorney, Agent or Firm: SCHWEGMAN LUNDBERG & WOESSNER, P.A.

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a continuation of prior application Ser. No. 17/484,667, filed on Sep. 24, 2021 and now issued as U.S. Pat. No. 11,922,163, which is incorporated by reference herein in its entirety.

BACKGROUND

(1) Traditionally, software has been self-contained and executed on one or more local machines. An enterprise desiring to use a software tool builds an on-premises computing system and executes a

software application to provide the tool on that computing system. The software application may be developed by the enterprise and/or purchased from a third party software provider. Users access the software tool directly from the computing system or remotely via a networked user computing device.

Description

BRIEF DESCRIPTION OF DRAWINGS

- (1) The present disclosure is illustrated by way of example and not limitation in the following figures.
- (2) FIG. 1 is a diagram showing one example of an environment for upgrading a legacy on-premises software application in a cloud environment.
- (3) FIG. 2 is a flowchart showing one example of a process flow that may be executed in the environment of FIG. 1 to determine upgrades available for a given application instance.
- (4) FIG. 3 is a flowchart showing one example of a process flow that may be executed in the environment of FIG. 1 to implement an upgrade of an application instance.
- (5) FIG. 4 is a flowchart showing one example of a process flow that may be executed in the environment of FIG. 1 to generate and deploy new application versions.
- (6) FIG. 5 is a flowchart showing one example of a process flow that may be executed in the environment of FIG. 1 to back-out a previously-generated application version.
- (7) FIG. 6 is a diagram showing one example of an environment for providing versions of a legacy on-premises software application in a cloud environment.
- (8) FIG. 7 is a flowchart showing one example of a process flow that may be executed in the environment of FIG. 6 to make a list of permissible versions for a user group.
- (9) FIG. 8 is a flowchart showing one example of a process flow that may be executed in the environment of FIG. 6 to build the list of permissible versions for a user group using the rollout strategy map and consumer context data.
- (10) FIG. 9 is a flowchart showing one example of a process flow that may be executed in the environment of FIG. 6 to determine if selector data from a rollout record matches consumer context data for a user group.
- (11) FIG. 10 is a flowchart showing one example of a process flow that may be executed in the environment of FIG. 6 to respond to a request for an indication of the permissible versions for a user group.
- (12) FIG. 11 is a flowchart showing one example of a process flow that may be executed in the environment of FIG. 6 to respond to a request for available upgrades to an application instance using the upgrade strategy map.
- (13) FIG. 12 is a flowchart showing one example of a process flow that may be executed in the environment of FIG. 6 to create a new instance of an application for a user group.
- (14) FIG. 13 is a flowchart showing one example of a process flow that may be executed in the environment of FIG. 6 to upgrade an instance of an application for a user group.
- (15) FIG. 14 is a flowchart showing one example of a process flow that may be executed in the environment of FIG. 6 to maintain the rollout strategy map to make an application version available to a particular user group.
- (16) FIG. 15 is a diagram showing one example of an arrangement for maintaining versions of a legacy on-premises software application in a cloud environment using a version metadata record.
- (17) FIG. 16 is a flowchart showing one example of a process flow that may be executed by in the arrangement of FIG. 15 to launch the application instance using instance metadata.
- (18) FIG. 17 is a flowchart showing one example of a process flow that may be executed in the arrangement of FIG. 15 to update the version metadata persistence.

(19) FIG. 18 is a process flow that may be executed in the arrangement of FIG. 15 to maintain an application instance.

(20) FIG. 19 is a flowchart showing one example of a process flow for updating the application instance in the example arrangement of FIG. 15.

(21) FIG. 20 is a block diagram showing one example of a software architecture for a computing device.

(22) FIG. 21 is a block diagram of a machine in the example form of a computer system within which instructions may be executed for causing the machine to perform any one or more of the methodologies discussed herein.

DETAILED DESCRIPTION

(23) An on-premises software application is typically executed by a computing system built by an enterprise on-site or on-premises. An on-premises application may be implemented to include a set of one or more executables, libraries, and the like, that are implemented by a set of executable files, libraries, and the like that are resident at the on-premises computing system. Users associated with the enterprise access the software application from the on-premises computing system.

(24) On-premises software applications are typically characterized by long release cycles. For example, changes to one functionality of the on-premises application may be implemented by changing one or more executables and/or libraries. The executables, libraries, and the like of the on-premises application, however, will typically implement multiple functionalities. This means that changes to one functionality of an on-premises application may require the re-testing of many different functionalities provided by the modified executable and/or library files. Also, when the on-premises application is a database management system, or other application interfacing with a data persistence, changes to the executable(s) and/or libraries for implementing the application may change the way that the application interfaces with the persistence. Accordingly, upgrading the on-premises application may include performing a data migration at the persistence (e.g., from one schema to a target schema). As a result, changes to an on-premises application may be made in batches. In this way, a batch of multiple changes can be tested, de-bugged, and deployed at the same time.

(25) In contrast to an on-premises implementation, some software applications are executed in a cloud environment. A cloud environment includes one or more data centers implementing one or more virtual and/or hardware servers. The cloud environment executes instances of cloud applications. For example, an instance of a cloud application may execute to provide a software tool to a group of users (e.g., a group of users associated with an entity that purchases access to the software tool). Each instance of the cloud application comprises a set of one or more virtual machines and/or containers executing at the cloud environment.

(26) Cloud-native applications are specifically developed to execute in a cloud environment. As a result, many cloud-native applications are designed according to a microservice architecture. In a microservice architecture, an instance of a software application is implemented by a collection of loosely-coupled microservices executing at the cloud environment. Each microservice may also include a single executable that executes in a separate virtual machine (VM) or container implemented by the cloud environment. In a microservice architecture, each microservice is programmed to perform a defined task or small set of tasks and interact with the other microservices in a defined way, for example, according to an application programming interface (API).

(27) The degree of independence between microservices in a microservice-architecture application simplifies the process of testing and upgrading. For example, due to the nature of microservice architecture, a change to one microservice may have a limited effect on other microservices and/or on data persistences used by the microservice-architecture application. Accordingly, many cloud-native microservice-architecture applications support testing and updating via a continuous or ongoing process, such as a continuous integration/continuous delivery (CI/CD) pipeline. In a

CI/CD pipeline, a developer can make changes to a subset of the microservices making up an application instance. The changes are implemented as a modified build of the application. The modified build is deployed to various integration and testing environments. Provided that the modified build passes tests at the various integration and testing environments, it may be deployed to one or more cloud environments as a new build. If the new build is found to be defective or to include bugs, it can be rolled back by reverting to a previous version of the build.

(28) In some examples, it is desirable to transition a software application designed for on-premises execution to a cloud environment. Because of the differences between a cloud environment and an on-premises computing system environment, such a transition may be non-trivial. In some examples, the executables, libraries, and the like of the on-premises application are configured to execute in one or more containers at the cloud environment. Orchestration services at the cloud environment may manage the containers executing at the cloud environment to provide instances of the application to groups of users associated with different customers.

(29) In some examples, a legacy on-premises application can be fully re-factored and re-coded to conform to a microservices architecture. Once such a re-factoring and re-coding occurs, the now-cloud-native application may be continuously upgraded using a CI/CD pipeline or similarly continuous technique. Oftentimes, however, it is not desirable to perform a full re-factoring and re-coding of a legacy on-premises application. For example, such a conversion may be time and labor-intensive. Instead of a full re-factoring and recoding, an on-premises application may be subject to a partial re-factoring. In some examples, the existing executable(s) and libraries of an on-premises application are just configured to execute in a cloud-environment-executed container managed by one or more cloud services.

(30) When an on-premises application is not fully-converted to a microservices architecture format, the resulting cloud implementation of the application may be less than optimally-suited for continuous updating. For example, container-implemented executables and/or libraries adapted from an on-premises implementation may not be as independent as a cloud-native applications. Accordingly, changes to one functionality implemented by an executable or library may entail more extensive testing of the other functionalities implemented by the same executable or library or other executables and/or libraries.

(31) Also, a cloud application with container-implemented executables and/or libraries may not be independent of the schema or schemas of the persistences used by the application. Accordingly, updating the application at the cloud may include performing a data migration at one or more cloud persistences to make those persistences consistent with the upgraded application. Such data migrations may be time and resource-intensive and, therefore, difficult to roll back in the event of a bug or other error.

(32) Another example problem with cloud-implemented legacy on-premises applications is that the testing of such applications may be more difficult than with on-premises implementations. For example, an on-premises application may be tested according to a process that involves providing various development versions (e.g., alpha, beta, etc.) to user groups for evaluation and/or early adoption. For example, testers or early adopters of an on-premises application may simply be provided with the relevant executables and/or libraries, which can be loaded on a local machine and/or network for testing. With a cloud implementation, however, making upcoming evaluation or early-adopter versions of an application available to stakeholder user groups is more challenging.

(33) Various examples described herein address these and other issues by utilizing application orchestration services and a product version map. A cloud environment may be arranged to execute instances of different versions of a software application for different groups of users. Each group of users may be associated with a stakeholder, such as a testing or quality assurance (QA) group at the entity providing the software application, a customer using the application, and the like.

(34) The product version map for a software application may include references to different available versions of the software application. The references to different available versions may

include identifiers and/or links to one or more container images for executing the referred-to containers, such as Docker® images provided by Docker, Inc. Application orchestration services may provide users with indications of the available versions of the application. A user may select a version of the application for that user and/or for a group of users including the user. In response, the application orchestration services may spin up an instance of the selected version of the application by initiating a container or containers described by the relevant container image. In this way, different users or groups of users may utilize different versions of the application. This may facilitate the testing of legacy on-premises applications that are implemented in the cloud environment. For example, a user group associated with QA or testing may spin up an evaluation version of the application for testing. A user group associated with a preferred customer entity may spin up and test a beta or other evaluation version of the application, and so on.

(35) Releases of the application can be managed by modifying the product version map. For example, when an evaluation version of the application is ready to launch, the product version map may be reconfigured to refer to the previous evaluation version as a new production version.

(36) Use of a product version map, however, may present its own challenges. For example, it may be desirable to program the orchestration services to be agnostic as to the relationship between different versions of the application. In this way, the orchestration services themselves may not need to be re-coded every time there is a change to the available versions. As a consequence of making the orchestration services agnostic as to the relationship between different versions, however, it may become possible for users or user groups to perform changes on their application instances that may not be advisable. For example, a user group may select a version of the application that is not compatible with a persistence associated with that user group and for which no data migration is defined. Also, for example, some available versions of an application may represent a downgrade relative to what the user group is already using.

(37) In various examples, these and other issues may be addressed utilizing an upgrade strategy map. An upgrade strategy map describes allowable upgrades from various source versions of the application to one or more allowable target versions of the application. The upgrade strategy map may comprise upgrade strategy records, where each upgrade strategy record lists a source version of the application and one or more target versions to which the application can be upgraded. When an upgrade is requested for a user or group of users, the orchestration service may retrieve an update strategy record corresponding to the current version of the application for that user group. The update strategy record may comprise an indication of a set of target versions of the application for the current application version. The set of target versions may indicate versions of the application that it is permissible for the group of users to upgrade to from the current version used by the group of users. The set of target versions may be provided to a user or users from the group of users, who may select a target version. In response, the application orchestration service may generate and execute one or more containers at the cloud environment to implement an instance of the selected target version of the application.

(38) FIG. 1 is a diagram showing one example of an arrangement **100** for upgrading a legacy on-premises software application in a cloud environment **102**. The cloud environment **102** executes one or more application instances **122, 124, 126, 128** and various services **112, 104** to provide an application, such as a database management application, to one or more user groups **114, 116, 118, 120**.

(39) The cloud environment **102** may include one or more computing devices, such as servers, for executing the application instances **122, 124, 126, 128** and services **112, 104**. In some examples, the cloud environment **102** is one of a number of cloud environments implemented by a cloud service provider at different geographic locations. For example, the cloud environment **102** may be implemented at a data center or data centers geographically near to the user groups **114, 116, 118, 120** so as to minimize network latencies to the user groups **114, 116, 118, 120**. Other cloud environments similar to the cloud environment **102** may be implemented at data centers

geographically near to other user groups to provide the application to those user groups.

(40) User groups **114, 116, 118, 120** indicate groups of one or more users who develop and/or use applications instances **122, 124, 126, 128** at the cloud environment **102**. In some examples, each user group **114, 116, 118, 120** is associated with a stakeholder. For example, some user groups **114, 116, 118, 120** are associated with stakeholders who are customer entities. A customer entity purchases or otherwise acquires access to application instances **122, 124, 126, 128**. A user group **114, 116, 118, 120** associated with a customer entity may include users who utilize one or more application instances **122, 124, 126, 128**, for example, use the functionality of the applications. Also, for example, some user groups **114, 116, 118, 120** are associated with internal stakeholders of an entity providing the software application. For example, some user groups **114, 116, 118, 120** may be associated with a development group for developing some or all of the application instances **122, 124, 126, 128**, a QA and/or testing user group for testing application instances **122, 124, 126, 128**, and/or the like. User groups **114, 116, 118, 120** may be associated with one or more application instances **122, 124, 126, 128**. For example, the user group **114** may be associated with application instance **122**. The user group **116** may be associated with the application instance **124**. The user group **118** may be associated with application instance **126**. The user group **120** may be associated with application instance **128**, and so on. In some examples, a user group **114, 116, 118, 120** may be associated with more than one instance of the application. For example, a user group **114, 116, 118, 120** may execute an instance of the application for a production implementation and another instance or instances of the application for a testing or evaluation use.

(41) The cloud environment **102** also implements persistences **130, 132, 134, 136**. For examples, when the application instances **122, 124, 126, 128** implement a database management system, the persistences **130, 132, 134, 136** may store data managed by the database management application. The persistences **130, 132, 134, 136** may be implemented using any suitable data storage device or devices such as, for example, one or more disks, one or more solid state drives, one or more random access memories, etc. In some examples, the software application implemented at the cloud environment **102** is a database management application for an in-memory database such as, for example, the HANA® and S/4 HANA® systems available from SAP SE of Walldorf, Germany.

(42) A persistence **130, 132, 134, 136** may be associated with a user group. In the example of FIG. 1, the persistence **130** is associated with user group **114**. The persistence **132** is associated with user group **116**. The persistence **134** is associated with the user group **118**; and the persistence **136** is associated with the user group **120** and so on. In some examples, a single user group **114, 116, 118, 120** may be associated with more than one persistence. For example, a user group **114, 116, 118, 120** may maintain one persistence for a production implementation of the software application and another persistence for a test or evaluation implementation of the software application.

(43) The application instances **122, 124, 126, 128** include a set of one or more executables, libraries, and/or other components executed within a single container or a limited set of containers implemented at the cloud environment. For example, the application instances **122, 124, 126, 128** may be instances of legacy on-premises applications executing within the cloud-implemented containers. In some examples, instances **122, 124, 126, 128** may be associated with respective persistences **130, 132, 134, 136** as shown. In some examples, containers for executing the various application instances **122, 124, 126, 128** are arranged as Kubernetes® clusters.

(44) The cloud environment **102** also executes one or more services **104, 112** for managing the application instances **122, 124, 126, 128**. The services **104, 112** may execute at the cloud environment **102** in one or more containers. A frontend service **112** may interface with the various user groups **114, 116, 118, 120**. For example, the frontend service **112** may provide users from various user groups **114, 116, 118, 120** with a user interface to allow the users to manage application instances **122, 124, 126, 128** that are executing and/or to launch new application instances. For example, the frontend service **112** may provide the various user groups **114, 116, 118, 120** with functionality to create, upgrade, delete, and/or otherwise manage application

instances **122, 124, 126, 128**. In some examples, the frontend service **112** also provides users from the various user groups **114, 116, 118, 120** with data about available target versions to which the user groups' own executing instances **122, 124, 126, 128** can be upgraded, as described herein.

(45) Application orchestration services **104** may manage the execution of the application instances **122, 124, 126, 128**. The application orchestration service **104** may be implemented as a single service and/or as a set of one or more subservices. In the example of FIG. 1, the application orchestration services **104** include an application operator service **105** and an application updater service **103**. In some examples, the application orchestration services **104** implement a container orchestration system, such as a Kubernetes® container orchestration system.

(46) The application operator service **105** may manage the execution of the one or more containers implementing an application instance **122, 124, 126, 128**. For example, the application operator service **105** may begin the execution of one or more containers for implementing an application instance **122, 124, 126, 128**, for example, from one or more container images associated with the one or more containers. Also, if one or more containers associated with an application instance **122, 124, 126, 128** crash, the application operator service **105** may spin up a replacement for the crashed container or containers. The application updater service **103** may manage application versions and update version metadata persistence **106** describing the application.

(47) Version metadata persistence **106** may include metadata describing various different versions of an application. The version metadata persistence **106** may include a product version map **108** and an upgrade strategy map **110**. The product version map **108** includes data describing available versions of the application. In some examples, the data describing a version may include release cycle data describing a maturity level of a version. For example, the release cycle data may indicate a level of testing and development to which the version has been subjected. Release cycle data may also indicate a frequency at which new versions are released. For example, a version may represent a monthly release, a bi-weekly release, or any other suitable release period. The product version map **108** may also include track data for one or more versions. The track data may indicate a status or upgrade track of the version and may be associated with a generally-available or production version. For example, the most recently-released quarterly version may be associated with a “latest” track. One or more previously-released versions may be associated with a “maintenance” track indicating that the version is not the latest but is still maintained.

(48) An example product version map **108** is given by EXAMPLE 1 below:

Example 1

(49) TABLE-US-00001 apiVersion: feature.example.com/v1alpha1 kind: ProductVersion metadata:
name: APPLICATION 1 spec: enablement: rolloutstrategy: staged-APP1-rollout
versions: - id: <some-unique-id> name: APP1 Cloud Edition 2019.38 releaseCycle: evaluation
image: some.container.registry/APP1-k8s:4.38.0 - id: <some-unique-id> name: APP1 Cloud
Edition 2019.36 releaseCycle: early-adoption image: some.container.registry/APP1-k8s:4.36.0
expiration-date: “2029-11-30” - id: <some-unique-id> name: APP1 Cloud Edition 2019.32
releaseCycle: generally-available-quarterly track: latest image: some.container.registry/APP1-
k8s:4.32.0 expiration-date: “2019-11-02” - id: <some-unique-id> name: APP1 Cloud Edition
2019.34 releaseCycle: generally-available-biweekly image: some.container.registry/APP1-
k8s:4.34.0 expiration-date: “2019-04-02” - id: <some-unique-id> name: APP1 Cloud Edition
2019.26 releaseCycle: generally-available-quarterly track: maintenance image:
some.container.registry/APP1-k8s:4.26.0 expiration-date: “2019-09-21”

(50) EXAMPLE 1 shows an example product version map **108** for an application called APPLICATION 1. The example product version map **108** shown in EXAMPLE 1 may represent the entirety of a product version map and/or may be a version map record for APPLICATION 1 retrieved from a product version map having product version map records for multiple different applications. The example product version map **108** indicates that APPLICATION 1 is subject to a staged rollout. The example product version map **108** also indicates five example versions called

APP1 Cloud Edition 2019.38, 2019.36, 2019.32, 2019.34, and 2019.26. In EXAMPLE 1, each version is further described by a unique identifier, a release cycle, and an expiration date. (51) The release cycle indicated for each version describes a maturity level of the version. For example, the version APP1 Cloud Edition 2019.38 is indicated to be an “evaluation” version. An evaluation version may be at a maturity level that is suitable for an early stage of testing and evaluation, such as may be performed by user groups **114, 116, 118, 120** that are internal QA or testing groups. The version APP1 Cloud Edition 2019.36 is indicated to be an “early-adoption” version. An early adoption version may be at a maturity level that is suitable for user groups **114, 116, 118, 120** who may desire to receive early or pre-release versions of the application, for example, for evaluation or testing. The versions APP1 Cloud Edition 2019.32 and 2019.26 are indicated to be “generally-available-quarterly” versions. This indicates versions at a maturity level that is tested and ready for general use. The “generally-available-quarterly” maturity level may indicate a release cycle of the application that includes a new generally-available version each quarter. The version APP1 Cloud Edition 2019.34 is indicated to be “generally-available-biweekly.” This indicates that the version is at a maturity level that is tested and ready for general use, although this version is released on a biweekly release schedule providing a new version every two weeks.

(52) The product version map of EXAMPLE 1 also indicates tracks for APP1 Cloud Edition 2019.32 and 2019.26, the two versions that have a maturity level indicating that they are generally available on a quarterly basis. For example, the upgrade track of version 2019.32 is “latest” while the upgrade track of version 2019.26 is labeled “maintenance.” For example, version 2019.26 may be a quarterly release from a previous quarter that is not the latest, but is still maintained.

(53) Version metadata persistence **106** may also include an upgrade strategy map **110**. The upgrade strategy map **110** describes allowable upgrades from various versions of the application. For example, the upgrade strategy map **110** may comprise upgrade strategy records, where each upgrade strategy record indicates a source version of the application and one or more target versions (if any) to which the source version may be upgraded. An example upgrade strategy map **110** is given by EXAMPLE 2 below:

Example 2

(54) TABLE-US-00002 apiVersion: feature.example.com/v1alpha1 kind: UpgradeStrategy
metadata: name: staged-APPLICATION 1-rollout-upgrades labels:
feature.example.com/productVersion: APP1 spec: upgrades: - source: selector:
releaseCycle: generally-available-quarterly track: maintenance targets: - name: patch selector:
releaseCycle: generally-available-quarterly track: maintenance - name: upgrade-to-latest-
quarterly selector: releaseCycle: generally-available-quarterly track: latest - name: upgrade-to-
biweekly selector: releaseCycle: generally-available-biweekly - name: upgrade-to-early-adoption
selector: releaseCycle: early-adoption - name: upgrade-to-evaluation selector: releaseCycle:
evaluation - source: selector: releaseCycle: generally-available-quarterly track: latest
targets: - name: patch selector: releaseCycle: generally-available-quarterly track: latest - name:
upgrade-to-biweekly selector: releaseCycle: generally-available-biweekly - name: upgrade-to-
early-adoption selector: releaseCycle: early-adoption - name: upgrade-to-evaluation selector:
releaseCycle: evaluation - source: selector: releaseCycle: generally-available-biweekly targets:
- name: patch selector: releaseCycle: generally-available-biweekly - name: switch-to-qrc
selector: releaseCycle: generally-available-quarterly - name: upgrade-to-early-adoption selector:
releaseCycle: early-adoption - name: upgrade-to-evaluation selector: releaseCycle: evaluation
- source: selector: releaseCycle: early-adoption targets: - name: patch selector:
releaseCycle: early-adoption - name: upgrade-to-evaluation selector: releaseCycle: evaluation
- source: selector: releaseCycle: evaluation targets: - name: patch selector: releaseCycle:
evaluation

(55) EXAMPLE 2 includes a number of example upgrade strategy records for the example

Application 1. Each upgrade strategy record indicates a source version and a set of one or more target versions to which the source version can be upgraded. In this example, source versions are indicated by maturity level given by the release cycle and upgrade track (if applicable). Also, in EXAMPLE 2, target versions are identified by a name, indicating the type of upgrade to that target version, and by the maturity level and upgrade track, if applicable, of the referenced version.

(56) Consider the first upgrade strategy record in EXAMPLE 2. The source version is identified by a release cycle “generally-available-quarterly” and a track, “maintenance.” As applied to the product version map **108** of EXAMPLE 1, this would refer to APP1 Cloud Edition 2019.26. A “patch” upgrade would upgrade to a target version that is the “generally-available-quarterly” version also on the “maintenance” track. An “upgrade-to-latest-quarterly” upgrade would upgrade to a target version that is the “generally-available-quarterly” version on the “latest” track. An “upgrade-to-biweekly” upgrade would change to the “generally-available-biweekly” version. An “upgrade-to-early-adoption” change would upgrade to the “early-adoption” version. An “upgrade-to-evaluation” change would upgrade to the “evaluation version.”

(57) In some examples, the application orchestration service **104** (e.g., via the application updater service **103** and/or application operator service **105**) reads the version metadata persistence **106** and prepares application instances **122, 124, 126, 128** for upgrading. For example, the application orchestration service **104** (e.g., via the application operator service **105**) may utilize the product version map **108** and upgrade strategy map **110** to determine the available upgrades for the various application instances **122, 124, 126, 128**. This may include determining the version and/or maturity level and upgrade track (if any) of an application instance **122, 124, 126, 128** and then determining, from the upgrade strategy map **110**, any upgrades that are available from that maturity level and upgrade track (if any).

(58) The application orchestration service **104** (e.g., application operator service **105**) may then consult the product version map **108** to determine the linked container images and version numbers associated with the target versions of each available upgrade. These may be written to a location accessible to the application orchestration service **104** (e.g., the application updater service **103** thereof), which may utilize the indicated target versions to respond to an upgrade request from a user group **114, 116, 118, 120**. When an upgrade request is received, the application orchestration service **104** (e.g., the application updater service **103**) may spin up a new version of the requested target version.

(59) In some examples, instead of pre-figuring the target versions from a particular application instance **122, 124, 126, 128**, the application orchestration service (e.g., the application operator service **105**) may consult the version metadata persistence **106** to determine available upgrades and target versions upon receiving a request for an upgrade from a user group **114, 116, 118, 120**. The arrangement **100** of FIG. 1 also includes a deployment service **138**. The deployment service **138** may execute at one or more computing devices, such as computing devices implemented by the cloud service provider that implements the cloud environment **102**. In some examples, the deployment service **138** is executed at the same computing devices that implement the environment **102**. The deployment service **138** may be programmed to launch the cloud environment **102**, for example, as part of a CI/CD pipeline. In some examples, the deployment service **138** may also be used to update the version metadata persistence **106**, for example, as described herein. For example, the CI/CD pipeline may be used to generate new versions, patches, etc. for the application instances **122, 124, 126, 128**. When the deployment service **138** deploys new version, patch, upgrade, or the like, it may also make appropriate upgrades to the version metadata persistence **106**.

(60) FIG. 2 is a flowchart showing one example of a process flow **200** that may be executed by the application orchestration service **104** (e.g., the application operator service **105**) to determine upgrades available for a given application instance **122, 124, 126, 128**. At operation **202**, the application orchestration service **104** receives an indication of an application instance **122, 124,**

126, 128. At operation **204**, the application orchestration service **104** determines the current version of the application instance **122, 124, 126, 128**. At operation **206**, the application orchestration service **104** determines the maturity level of the application instance **122, 124, 126, 128**. This may include, for example, using a version name and/or number to pull a product version record for the version from the product version map **108**. The product version record may provide release cycle and/or upgrade track data for the application instance **122, 124, 126, 128**.

(61) At operation **208**, the application orchestration service **104** retrieves the current upgrade strategy record for the application instance **122, 124, 126, 128**. This may involve, for example, identifying the upgrade strategy record at the upgrade strategy map **110** that lists the current version and/or maturity level of the application instance **122, 124, 126, 128** as the source version. At operation **210**, the application orchestration service **104** determines upgrades that are available to the application instance **122, 124, 126, 128**. If any upgrades are available at operation **212**, the application orchestration service **104**, at operation **216**, reads the version records for the indicated target versions from the product version map **108**. For example, as illustrated with EXAMPLES 1 and 2 above, the upgrade strategy record for a source version may indicate the maturity level of target versions. The application orchestration service **104** may utilize the product version map to find the version number, docker image(s), etc., associated with the indicated maturity level. The version records may be added to a set of target versions at operation **218**.

(62) At operation **220**, the application orchestration service **104** annotates the application instance **122, 124, 126, 128** to include the available upgrades including, for example, the target versions. The annotation can be made at any suitable location including, for example, at a persistence available to the application orchestration service **104** and the frontend service **112**, and/or at a metadata store associated with the application instance **122, 124, 126, 128**.

(63) Referring back to operations **210** and **212**, if the retrieved upgrade strategy record does not indicate any available upgrades for the application instance **122, 124, 126, 128**, the application orchestration service **104** may create an empty list of available upgrades at operation **214** to indicate that no upgrades are available. Optionally, the empty list of available upgrades is annotated at operation **220** as described herein. In some examples, rather than creating an annotation of an application instance **122, 124, 126, 128**, the application orchestration service **104** may execute the process flow **200** in response to a request for available upgrades from the frontend service **112**. For example, instead of annotating the available upgrades at operation **220**, the application orchestration service **104** may send a response to the frontend service **112** indicating the available upgrades.

(64) FIG. 3 is a flowchart showing one example of a process flow **300** that may be executed in the arrangement **100** to implement an upgrade of an application instance. The process flow **300** includes three columns **301, 303, 305** with each column including operations that may be executed by a different component of the arrangement **100**. The column **301** includes operations that can be performed by a user group **114, 116, 118, 120** or, for example, a single user authorized to make changes on behalf of the user group **114, 116, 118, 120**. The column **303** includes operations executed by the frontend service **112**. The column **305** includes operations performed by the application orchestration service **104**, such as by the application updater service **103**.

(65) At operation **302**, the user group **114, 116, 118, 120** opens an application dashboard. The application dashboard may be any suitable user interface element provided to the user group **114, 116, 118, 120** by the frontend service **112**. At operation **304**, the frontend service **112** reads available application versions. The available application versions are based on the version of the application instance **122, 124, 126, 128** executing for the corresponding user group **114, 116, 118, 120**. The available application versions, for example, may have been previously annotated by the application orchestration service **104** as described with respect to the process flow **200** of FIG. 2. Also, in some examples, the frontend service **112** directs a request for available versions to the application orchestration service **104**, which may respond by providing the requested list of

available versions.

(66) At operation **306**, the frontend service **112** notifies the user group of the available new versions. This may include, for example, generating a list of the available new versions that is provided to the user group **114, 116, 118, 120** via the dashboard UI element. At operation **308**, the user group **114, 116, 118, 120** chooses a patch or upgrade. For example, a patch may be a change to an application version that has the same maturity level as the current version of the application instance **122, 124, 126, 128** (e.g., the release cycle and/or upgrade track). An upgrade may be a change to an application version having a different maturity level than the current version of the application instance **122, 124, 126, 128** (e.g., a different release cycle and/or upgrade track). It will be appreciated that it may not be the case that both a patch and an upgrade are available in every circumstance. Sometimes, only one or more patches will be available. Sometimes only one or more upgrades will be available. In some examples, neither a patch nor an upgrade will be available.

(67) If the user group **114, 116, 118, 120** selects a patch, then the user group **114, 116, 118, 120** selects a target version in the same maturity level as the source version or current version at operation **312**. If the user group **114, 116, 118, 120** selects an upgrade, then, at **314**, the user group **114, 116, 118, 120** chooses a version in a different maturity level than the current version of the application. At operation **316**, the selected update or patch is provided to the frontend service **112**. At operation **318**, the frontend service **112** collects parameters for the selected update and calls the application orchestration service **104** (e.g., the application updater service **103**). The collected parameters may include, for example, an indication of the user group **114, 116, 118, 120** making the request, an indication of the current instance **122, 124, 126, 128** of the application for that user group **114, 116, 118, 120** including, for example, a version of the instance, and an indication of the selected target version.

(68) The application orchestration service **104** (e.g., the application updater service **103**), at operation **320**, retrieves one or more container images for launching the target version referred to at the product version map **108**. The one or more container images may be Docker® images or other suitable images describing containers to be executed at the cloud environment **102** to implement the selected target version. Each container may include one or more executables and/or libraries that execute within the container to implement an instance of the application version.

(69) In some examples, a selected target version may be associated with more than one container image. The orchestration service **104** may retrieve a container image or images based on selections made by the user group **114, 116, 118, 120**. For example, a selected target version may be executed with different containers or different combinations of containers based on the application features desired by the relevant user group **114, 116, 118, 120**.

(70) At operation **322**, the application orchestration service **104** (e.g., the application updater service **103**) upgrades the application instance. Upgrading the application instance may include, for example, spinning up one or more containers based on the retrieved container image or images to generate a new application instance for the requesting user group **114, 116, 118, 120**. In some examples, upgrading the application instance may also include executing a data migration at a persistence **130, 132, 134, 136** of the requesting user group **114, 116, 118, 120**. For example, the persistence **130, 132, 134, 136** or a portion thereof may be modified from a prior schema to a target schema used by the target version. In some examples, upgrading the application instance may also include stopping the execution of a container associated with a previous instance of the application used by the requesting user group **114, 116, 118, 120**.

(71) At operation **324**, the frontend service reports to the user group **114, 116, 118, 120** that the new instance of the application is ready for use. The user group **114, 116, 118, 120** may commence to use the new instance of the application at operation **326**.

(72) FIG. 4 is a flowchart showing one example of a process flow **400** that may be executed in the environment of FIG. 1 to generate and deploy new application versions. The process flow **400** includes three columns. A first column **401** shows operations that may be performed by an

application development team. The application development team may be a user group such as, for example, a user group **114, 116, 118, 120**, made up of developers, testing users, and the like. Column **403** shows operations that may be performed by a deployment service, such as the deployment service **138**. The column **405** shows operations that are executed by the application orchestration service **104** and/or a subservice thereof.

(73) At operation **402**, the application development team produces a new version of the application. The new version of the application may be arranged to execute at the cloud environment **102** in one or more containers, as described herein. At operation **404**, the application development team performs and/or supervises testing of the new version of the application. At operation **408**, the deployment service updates the upgrade strategy map at the cloud environment **102** in view of the new version. In some examples, this includes including the new version as an allowable target version for one or more source versions. While operations **402, 404, and 408** are being executed, the application orchestration service **104** may continue, at operation **406**, to perform upgrades to application instances **122, 124, 126, 128** using first version (V1) of the upgrade strategy map **110**. After the upgrade strategy map **110** is updated at operation **408**, the orchestration service **104** may, at operation **414**, begin to use a second version (V2) of the upgrade strategy map **110**, the second version being created by the update at operation **408**.

(74) In some examples, the changeover from V1 to V2 is transparent to the application orchestration service **104**. For example, the application orchestration service **104** may operate based on the version metadata persistence **106** installed at the cloud environment **102**. When the version metadata persistence **106** is modified, such as the change to the upgrade strategy map at operation **408**, the application orchestration service **104** may continue to operate, albeit based on the upgraded version metadata.

(75) At operation **410**, the application development team produces a second new version of the application. At operation **412**, the application development team performs and/or supervises testing of the second new version of the application. At operation **416**, the deployment service updates the upgrade strategy map **110** at the cloud environment **102** in view of the second new version. In some examples, this includes including the second new version as an allowable target version for one or more source versions. While operations **410, 412, and 414** are being executed, the application orchestration service **104** may continue to perform upgrades to application instances **122, 124, 126, 128** using V2 of the upgrade strategy map **110**. After the upgrade strategy map is updated at operation **416**, the orchestration service **104** may, at operation **418**, begin to use a third version (V3) of the upgrade strategy map **110**, the third version being created by the update at operation **416**.

(76) At operation **420**, the application development team produces a third new version of the application. At operation **422**, the application development team performs and/or supervises testing of the third new version of the application. At operation **424**, the deployment service updates the upgrade strategy map **110** at the cloud environment **102** in view of the third new version. In some examples, this includes including the third new version as an allowable target version for one or more source versions. While operations **420, 422, and 424** are being executed, the application orchestration service **104** may continue to perform upgrades to application instances **122, 124, 126, 128** using V2 of the upgrade strategy map **110**. After the upgrade strategy map is updated at operation **416**, the orchestration service **104** may begin to use a fourth version (V4) of the upgrade strategy map **110**, the fourth version being created by the update at operation **424**. This process may continue as additional versions are produced by the application development team.

(77) FIG. 5 is a flowchart showing one example of a process flow **500** that may be executed in the environment of FIG. 1 to back-out a previously-generated application version. A first column **501** shows operations that may be performed by an application development team such as, for example, a user group **114, 116, 118, 120** made up of developers, testing users, and the like. Column **503** shows operations that may be performed by a deployment service, such as the deployment service

138. The column **505** shows operations that are executed by that application orchestration service **104** and/or a subservice thereof.

(78) At operation **502**, the application development team produces a new version of the application. The new version of the application may be arranged to execute at the cloud environment **102** in one or more containers, as described herein. At operation **504**, the application development team performs and/or supervises testing of the new version of the application. At operation **508**, the deployment service updates the upgrade strategy map at the cloud environment **102** in view of the new version. In some examples, this includes including the new version as an allowable target version for one or more source versions. While operations **502**, **504**, and **508** are being executed, the application orchestration service **104** may continue, at operation **506**, to perform upgrades to application instances **122**, **124**, **126**, **128** using first version (V1) of the upgrade strategy map **110**. After the upgrade strategy map is updated at operation **508**, the orchestration service **104** may, at operation **510**, begin to use a second version (V2) of the upgrade strategy map **110**, the second version being created by the update at operation **508**.

(79) At operation **512**, the application development team detects that the new version of the application is erroneous. For example, the new version may have a QA issue, such as a reported error or other problem with the execution of or other aspect of the new version. In response, at operation **516**, the buggy upgrade path to the new version is backed out of the version metadata persistence **106** at the cloud environment **102**, for example, using the deployment service **138**. This may include, for example, deleting one or more references to the new version of the application as allowable target versions at the upgrade strategy map **110** at the environment **102**. The result of operation **516** may be a third version of the upgrade strategy map **110** at the environment **102** (V3) that does not include a reference to the buggy new version as a permitted target version. While operations **512** and **516** are being executed, the application orchestration service **104** may continue, at operation **514**, to use the second version of the upgrade strategy map **110** for generating allowable upgrades and/or performing the upgrades. After the execution of operation **516**, the application orchestration service **104** may begin, at operation **518**, to use the third version of the upgrade strategy map **110**.

(80) In some examples, it may be desirable to selectively expose some versions of an application to different user groups. Consider an example user group that includes users who are part of a QA department within the entity providing the application. It may be desirable to expose alpha, beta, and/or other early testing versions of an application to the example user group and not to other user groups, such as user groups associated with customer entities. Consider another example user group associated with a premium customer of the entity generating the application. It may be desirable to expose beta and/or other evaluation versions of the application to this example user group without exposing those versions to all customers.

(81) Selectively providing different versions of a legacy on-premises application to different users from a cloud environment may present challenges. For example, as described herein, it may be desirable for application orchestration services at the cloud environment to be independent of application versions so that application versions can be changed without requiring changes to the orchestration services. Similarly, it may be desirable to modify the versions of a legacy on-premises application available to different user groups without making changes to the orchestration services themselves.

(82) Various examples address these and other challenges utilizing a version catalog and rollout strategy map. The rollout strategy map comprises rollout records, where each rollout record describes an application version and selector data. The selector data describes user groups that are authorized to use the corresponding application version. The version catalog may include a reference to the rollout strategy map.

(83) FIG. **6** is a diagram showing one example of an arrangement **600** for providing versions of a legacy on-premises software application in a cloud environment **602**. The cloud environment **602**

executes one or more application instances **622, 624, 626, 628** and various services **612, 604** to provide an application, such as a database management application, to one or more user groups **614, 616, 618, 620**.

(84) The cloud environment **602** may include one or more computing devices, such as servers, for executing the application instances **622, 624, 626, 628** and services **612, 604**. In some examples, the cloud environment **602** is one of a number of cloud environments implemented by a cloud service provider at different geographic locations. For example, the cloud environment **602** may be implemented at a data center or data centers geographically near to the user groups **614, 616, 618, 620** so as to minimize network latencies to the user groups **614, 616, 618, 620**. Other cloud environments similar to the cloud environment **602** may be implemented at data centers geographically near to other user groups to provide the application to those user groups.

(85) User groups **614, 616, 618, 620** indicate groups of one or more users who use applications instances **622, 624, 626, 628** at the cloud environment **602**. Each user group **614, 616, 618, 620**, in some examples, is associated with a stakeholder. Stakeholders may be internal stakeholders of an entity providing the software application such as, for example, QA or testing groups. Stakeholders may also include customer entities who purchase access to the software application.

(86) User groups **614, 616, 618, 620** may be associated with one or more application instances **622, 624, 626, 628**. For example, the user group **614** may be associated with application instance **622**. The user group **616** may be associated with the application instance **624**. The user group **618** may be associated with application instance **626**. The user group **620** may be associated with application instance **628**, and so on. In some examples, a user group **614, 616, 618, 620** may be associated with more than one instance of the application. For example, a user group **614, 616, 618, 620** may execute an instance of the application for a production implementation and another instance or instances of the application for a testing or evaluation use.

(87) The cloud environment **602** also implements persistences **630, 632, 634, 636**. For examples, when the application instances **622, 624, 626, 628** implement a database management system, the persistences **630, 632, 634, 636** may store data managed by the database management application. The persistences **630, 632, 634, 636** may be implemented using any suitable data storage device or devices such as, for example, one or more disks, one or more solid state drives, one or more random access memories, etc. In some examples, the software application implemented at the cloud environment **602** is a database management application for an in-memory database such as, for example, the HANA® and S/4 HANA® systems available from SAP SE of Walldorf, Germany.

(88) A persistence **630, 632, 634, 636** may be associated with a user group. In the example of FIG. 6, the persistence **630** is associated with user group **614**. The persistence **632** is associated with user group **616**. The persistence **634** is associated with the user group **618**; and the persistence **636** is associated with the user group **620** and so on. In some examples, a single user group **614, 616, 618, 620** may be associated with more than one persistence. For example, a user group **614, 616, 618, 620** may maintain one persistence for a production implementation of the software application and another persistence for a test or evaluation implementation of the software application.

(89) The software application provided by the cloud environment **602** of FIG. 6 may be a legacy on-premises application converted for provision by the cloud environment **602**. For example, the application may execute as various instances **622, 624, 626, 628**. Each application instance may comprise one or more containers including executables, libraries, and/or other components for providing the software application. In some examples, instances **622, 624, 626, 628** may be associated with respective persistences **630, 632, 634, 636** as shown. In some examples, containers for executing the various application instances **622, 624, 626, 628** are arranged as Kubernetes® clusters.

(90) The cloud environment **602** also executes one or more services **604, 612** for managing execution of the application instances **622, 624, 626, 628**. The services **604, 612** may execute at the cloud environment **602** in one or more containers. A frontend service **612** may interface with the

various user groups **614, 616, 618, 620**. For example, the frontend service **612** may provide users from various user groups **614, 616, 618, 620** with a user interface to allow the users to access application instances **622, 624, 626, 628** that are executing and/or to launch new application instances. In some examples, the frontend service **612** provides users from the various user groups **614, 616, 618, 620** with data about available target versions to which the user groups' own executing instances **622, 624, 626, 628** can be upgraded, as described herein.

(91) Application orchestration services **604** may manage the execution of the application instances **622, 624, 626, 628**. The application orchestration service **604** may be implemented as a single service and/or as a set of one or more subservices. In the example of FIG. 6, the application orchestration services **604** include an application operator service **605** and an application updater service **603**.

(92) The application operator service **605** may manage the execution of the one or more containers implementing an application instance **622, 624, 626, 628**. For example, the application operator service **605** may begin the execution of one or more containers for implementing an application instance **622, 624, 626, 628**, for example, from one or more container images associated with the one or more containers. Also, if one or more containers associated with an application instance **622, 624, 626, 628** crash, the application operator service **605** may spin up a replacement for the crashed container or containers. The application updater service **603** may manage application versions.

(93) In the example of FIG. 6, the version metadata persistence **606** includes a product version map **608**, an upgrade strategy map **610** and a rollout strategy map **609**. The product version map **608** and upgrade strategy map **610** may be similar to the product version map **108** and upgrade strategy map **110** described herein with respect to FIG. 1. The rollout strategy map **609** may include rollout records, with each rollout record describing a version of the application and selector data, where the selector data describes user group context properties describing user groups that are permitted to use the indicated application version.

(94) The arrangement **600** also includes a version catalog service **607**. The version catalog service **607** may be used by the application orchestration service **104** and/or the frontend service **612** to access the version metadata persistence **606**. For example, the version catalog service **607** may utilize the product version map **608**, rollout strategy map **609**, and/or upgrade strategy map **110** to generate lists of permissible versions for some or all of the user groups **614, 616, 618, 620**. A list of permissible versions for a user group **614, 616, 618, 620** indicates versions of the application that are both available and that the corresponding user group **614, 616, 618, 620** is authorized to use. The version catalog service **607** may generate and/or update a list of permissible versions for a user group **614, 616, 618, 620**, for example, upon a request from a user group **614, 616, 618, 620** to initiate and/or upgrade an application instances **622, 624, 626, 628** and/or periodically.

(95) In some examples, the application orchestration service **604** may request that the version catalog service **607** utilize the rollout strategy map **609** in conjunction with the upgrade strategy map **610**. For example, the list of permissible versions for a user group **614, 616, 618, 620** at the version catalog service **607** may include versions that the user group **614, 616, 618, 620** is permitted to use, as indicated by the rollout strategy map **609** and that are permissible upgrades from the current version of an application instance **622, 624, 626, 628** for the user group **614, 616, 618, 620**.

(96) An example rollout strategy map **609** is provided by EXAMPLE 3 below. The example rollout strategy map **609** provided by EXAMPLE 3 below may be used in conjunction with the example product version map given by EXAMPLE 1 herein and/or with the example upgrade strategy map **610** given by EXAMPLE 2 herein.

Example 3

(97) TABLE-US-00003 apiVersion: feature.example.com/v1alpha1 kind: RolloutStrategy metadata: name: staged-APP1-rollout labels: feature.example.com/productVersion: APP1 spec:

default: versionSelector: releaseCycle: generally-available-quarterly track: qrc4 rollouts:

- name: QRC selector: releaseCycle: generally-available-quarterly track: qrc3

versionSelector: releaseCycle: generally-available-quarterly track: qrc3 - name: biweekly selector: releaseCycle: biweekly versionSelector: releaseCycle: generally-available-biweekly (98) The name of the rollout strategy map example of EXAMPLE 3 is “staged-APP1-rollout.” For example, this rollout strategy map example was referred to in the product version map of EXAMPLE 1 above. The example rollout strategy map of EXAMPLE 3 includes three rollout records. A first default rollout record called “default” indicates a default version of APPLICATION 1 that is to be available, which is the “generally-available-quarterly” version on track quarterly release cycle (QRC 3). The first rollout record does not indicate any selector data, for example, because the indicated version is the default version that is available.

(99) Other example rollout records indicated by EXAMPLE 3 include a rollout record named QRC and a rollout record named “biweekly.” The QRC rollout record is associated with selector data indicating a release cycle—“generally-available-quarterly” and an upgrade track “qrc 3.” The record indicates an application version corresponding to the “generally-available-quarterly” release cycle and the qrc3 track, which can be tied to a particular application version at the product version map 608. This rollout record indicates user groups 614, 616, 618, 620 having context properties indicating that the user group 614, 616, 618, 620 is on a “generally-available-quarterly” release cycle and the upgrade track “qrc 3” is authorized to use the indicated version.

(100) The “biweekly” rollout record indicates an application version corresponding to the “generally-available-biweekly” release cycle and selector data indicating the “biweekly” release cycle. Accordingly, user groups 614, 616, 618, 620 having context properties indicating that the user group 614, 616, 618, 620 is on a “biweekly” release cycle are authorized to access the indicated version.

(101) The arrangement 100 of FIG. 6 also includes a deployment service 638. The deployment service 638 may execute at one or more computing devices, such as computing devices implemented by the cloud service provider that implements the cloud environment 602. In some examples, the deployment service 638 is executed at the same computing devices that implement the environment 602. The deployment service 638 may be programmed to launch the cloud environment 602, for example, as part of a CI/CD pipeline. In some examples, the deployment service 638 may also be used to update the version metadata persistence 606, for example, as described herein. For example, the CI/CD pipeline may be used to generate new versions, patches, etc. for the application instances 622, 624, 626, 628. When the deployment service 638 deploys new version, patch, upgrade, or the like, it may also make appropriate upgrades to the version metadata persistence 606.

(102) FIG. 7 is a flowchart showing one example of a process flow 700 that may be executed in the arrangement 600 by version catalog service 607 to make a list of permissible versions for a user group 614, 616, 618. The process flow 700 may be executed in response to a request from a user group 614, 616, 618, 620 (e.g., via the frontend service 612, in response to a request from the application orchestration service 604, periodically, and/or in response to any other suitable criteria.

(103) At operation 702, the version catalog service 607 accesses the rollout strategy map 609. At operation 704, the version catalog service 607 accesses consumer context data for a user group 614, 616, 618, 620. The context data comprises context properties describing the user group 614, 616, 618, 620. The consumer context data can be received from a user associated with the user group 614, 616, 618, 620 as part of a request to initiate and/or upgrade an application instance 622, 624, 626, 628. Also, in some examples, consumer context data for different user groups 614, 616, 618, 620 is archived. The version catalog service 607 may retrieve the consumer context data for a user group 614, 616, 618, 620 from an archive.

(104) Example context properties that may be stored with the consumer context data for a user group 614, 616, 618, 620 include a name of the application provided by the cloud environment

602, a customer account identifier, a customer sub-account identifier, an identifier of the instance **622**, **624**, **626**, **628** of the application currenting executing for the user group **614**, **616**, **618**, **620**, an indication of a release cycle for the user group **614**, **616**, **618**, **620**, an indication of an upgrade track for the user group **614**, **616**, **618**, **620**, and the like.

(105) At operation **706**, the version catalog service **607** may build a list of versions of the application that are permissible for the user group **614**, **616**, **618**, **620** in view of the rollout strategy map **609** and product version map **608**. The list of permissible versions may be returned at operation **708**. Returning the list of permissible versions may include providing the list to the frontend service **612** or other requestor.

(106) FIG. **8** is a flowchart showing one example of a process flow **800** that may be executed in the arrangement **600** to build the list of permissible versions for a user group **614**, **616**, **618**, **620** using the rollout strategy map **609** and consumer context data. The process flow **800** shows one example way that the version catalog service **607** can execute the operation **706** of the process flow **700**.

(107) At operation **802**, the version catalog service **607** accesses a rollout record from the rollout strategy map **609**. At operation **804**, the version catalog service **607** determines if the consumer context data matches the selector data of the rollout record. In some examples, if the rollout record is a default rollout record for a default version of the application, then operation **804** may be considered to match all consumer context data. If the rollout record is not for a default version of the application, then the operation **804** may include comparing the consumer context properties of the user group **614**, **616**, **618**, **620** to the selector data of the rollout record.

(108) If, at operation **804**, the rollout record selector data matches the consumer context data, the version catalog service **607**, at operation **806**, accesses the application version indicated by the rollout record and, at operation **808**, adds the indicated application version to the list of permissible versions for the user group **614**, **616**, **618**, **620**. If, at operation **804**, the rollout record does not match the consumer context data, then the version catalog service **607** may proceed to operation **810**.

(109) After adding the application version to the list of permissible versions at operation **808** or determining, at operation **804**, that the rollout record selector data does not match the consumer context data, the version catalog service **607**, at operation **810**, determines if the rollout strategy map **609** includes any additional rollout records. If there are more rollout records, the version catalog service **607** accesses the next rollout record at operation **802** and proceeds as described herein. If there are no additional rollout records, the version catalog service **607** returns the current list of permissible versions for the user group **614**, **616**, **618**, **620** at operation **812**.

(110) FIG. **9** is a flowchart showing one example of a process flow **900** that may be executed in the arrangement **600** to determine if selector data from a rollout record matches consumer context data for a user group **614**. For example, the process flow **900** shows one example way that the version catalog service **607** can execute the operation **804** of the process flow **900**.

(111) At operation **902**, the version catalog service **607** accesses a first selector property from the selector data. At operation **904**, the version catalog service **607** determines if the first selector property accessed at operation **902** has a matching property at the consumer context data describing the user group **614**, **616**, **618**, **620**. For example, if the rollout selector property is a release cycle, the operation **904** may include determining if the consumer context data includes an indication of a release cycle. If there is no corresponding property in the consumer context data, the version catalog service **607** returns no match between the user group **614**, **616**, **618**, **620** and considered rollout record at operation **910**.

(112) If there is a corresponding property in the consumer context data at operation **904**, the version catalog service **607** determines at operation **908** whether the value of the corresponding property of the consumer context data has the same value as the first consumer context property. For example, if the first rollout selector property indicates a bi-weekly release cycle, the operation **908** may include determining whether the consumer context data for the corresponding user group **614**, **616**,

618, 620 also indicates the same the bi-weekly release cycle. If the value of the corresponding property of the consumer context data is not the same value as the first selector property, the version catalog service **607** returns no match between the user group **614, 616, 618, 620** and considered rollout record at operation **910**.

(113) If, at operation **908**, the value of the corresponding consumer context data property is the same value as the first selector property, the version catalog service **607** determines at operation **912** whether there are any additional selector properties are indicated by the rollout selector data. If there are additional selector properties the version catalog service **607** accesses the next selector property at operation **902** and continues as described.

(114) If all selector properties from the rollout record have been considered, the version catalog service **607** returns an indication that there is a match between the user group **614, 616, 618, 620** and considered rollout record at operation **914**.

(115) The example process flow **900** implements a logic that returns a match between the consumer context data for a user group **614** and selector data from a rollout record when all properties that are included in both the rollout selector and the consumer context data have equivalent values and all properties in the rollout selector have equivalent, matching properties in the consumer context data. It will be appreciated, however, that alternative arrangements can also be used. In some examples, consumer context data may be considered to match selector data from a rollout record when each consumer context property has a matching property in the rollout selector data having an equivalent value.

(116) FIG. **10** is a flowchart showing one example of a process flow **1000** that may be executed in the arrangement **600** to respond to a request for an indication of the permissible versions for a user group **614, 616, 618, 620**. For example, the process flow **1000** may be executed when a user from a user group **614, 616, 618, 620** requests that an application instance **622, 624, 626, 628** be initiated and/or upgraded.

(117) At operation **1002**, the version catalog service **607** receives the request for a list of permissible application versions for a user group **614, 616, 618, 620**. For example, the request may be received from a user associated with a user group **614, 616, 618, 620**, from the frontend service **612**, or any other suitable source. At operation **1004**, the version catalog service **607** accesses the product version map **608** and rollout strategy map **609** associated with the application. At operation **1006**, the version catalog service **607** accesses consumer context data for the user group **614, 616, 618, 620**. The consumer context data may be included with the request and/or accessed from another source.

(118) At operation **1008**, the version catalog service **607** builds the requested list of permissible versions for the user group **614, 616, 618, 620**. For example, the operation **1008** may be performed as described here with respect to FIGS. **7-9**. At operation **1010**, the version catalog service **607** returns the list of permissible versions generated at operation **1008**. For example, the list of permissible versions may be provided to at least one user of the indicated user group **614, 616, 618, 620**. That user may select a version of the application from the list, causing an application instance **622, 624, 626, 628** associated with the user group **614, 616, 618, 620** to be initiated and/or upgraded as described herein.

(119) FIG. **11** is a flowchart showing one example of a process flow **1100** that may be executed in the arrangement **600** to respond to a request for available upgrades to an application instance **622, 624, 626, 628** using the upgrade strategy map **610**. For example, the process flow **1000** may be executed when a user from a user group **614, 616, 618, 620** requests that an application instance **622, 624, 626, 628** be upgraded.

(120) At operation **1102**, the version catalog service **607** receives a request for a list of available upgrades for a particular user group **614, 616, 618, 620**. At operation **1104**, the version catalog service **607** accesses the product version map **608**, rollout strategy map **609**, and upgrade strategy map **610** for the relevant application. At operation **1106**, the version catalog service **607** accesses

consumer context data for the user group **614, 616, 618, 620**. The consumer context data may be included with the request and/or accessed from another source.

(121) At operation **1108**, the version catalog service **607** accesses the version of the application instance **622, 624, 626, 628** currently executing for the user group **614, 616, 618, 620**. At operation **1110**, the version catalog service **607** builds a set of target versions to which the currently executing instance can be upgraded using the upgrade strategy map **610**. This can be performed, for example, as described herein with respect to FIG. 2. At operation **1112**, the version catalog service **607** builds a list of permissible versions for the user group **614, 616, 618, 620**, for example, as described herein with respect to FIGS. 7-9.

(122) At operation **1114**, the version catalog service **607** filters the list of permissible versions and the set of target versions to generate a filtered list of permissible versions. For example, the filtered list of permissible versions may include versions of the application that were both on the set of target versions and on the initial list of permissible versions generated at operation **1112**. At operation **1116**, the version catalog service **607** returns the filtered list of permissible versions. For example, the list of permissible versions may be returned to a user of the relevant user group **614, 616, 618, 620**, to the frontend service **612**, or to another suitable component.

(123) FIG. 12 is a flowchart showing one example of a process flow **1200** that may be executed in the arrangement **600** to create a new instance **622, 624, 626, 628** of an application for a user group **614, 616, 618, 620**. The process flow **1200** includes three columns **1201, 1203, 1205** with each column including operations that may be executed by a different component of the arrangement **600**. The column **1201** includes operations that can be performed by a user group **614, 616, 618, 620** or, for example, a single user authorized to make changes on behalf of the user group **614, 616, 618, 620**. The column **1203** includes operations executed by the frontend service **612**. The column **1205** includes operations performed by the application orchestration service **604**, such as by the application operator service **605**.

(124) At operation **1202**, the user group **614, 616, 618, 620** begins the provisioning process, for example, by opening and/or otherwise requesting an application instance via a dashboard or other suitable user interface provided by the frontend service **612**. At operation **1204**, the frontend service **612** builds consumer context data for the user group **614, 616, 618, 620**. This may include, for example, querying the user or users from the user group **614, 616, 618, 620** who opened the dashboard at operation **1202** and/or accessing data about the user and/or user group **614, 616, 618, 620**. At operation **1206**, the frontend service **612** requests the list of permissible versions for the user group **614, 616, 618, 620**. For example, the request may be directed to the version catalog service **607**. The version catalog service **607** may access a previously-generated list of permissible applications for the user group **614, 616, 618, 620** and/or may generate or request that the list of permissible versions be generated using the consumer context data generated at operation **1204**.

(125) The frontend service **612** receives the list of permissible versions at operation **1208** and provides the list to the user group at operation **1210**. For example, the frontend service **612** may provide the list of permissible versions via a user interface provided to a user or users of the user group **614, 616, 618, 620**. The user or users from the user group **614, 616, 618, 620** select a version for provisioning at operation **1212**.

(126) At operation **1216**, the frontend service **612** collects parameters for the selected version and calls the application orchestration service **604** (e.g., the application operator service **605**). The collected parameters may include, for example, an indication of the user group **614, 616, 618, 620** making the request and an indication of the selected version from the list of permissible versions.

(127) The application orchestration service **604** (e.g., the application operator service **605**), at operation **1218**, retrieves the container image of the target version referred to at the product version map **608** and, at operation **1220**, creates an application instances **622, 624, 626, 628** of the requested version for the user group **614, 616, 618, 620**. Creating the instance may include, for example, spinning up one or more containers to execute a new application instance for the

requesting user group **614, 616, 618, 620**.

(128) At operation **1222**, the frontend service **612** reports to the user group **614, 616, 618, 620** that the new instance of the application is ready for use. The user group **614, 616, 618, 620** may commence to use the new instance of the application at operation **1224**.

(129) FIG. **13** is a flowchart showing one example of a process flow **1300** that may be executed in the arrangement **600** to upgrade an instance **622, 624, 626, 628** of an application for a user group **614, 616, 618, 620**. The process flow **1300** includes three columns **1301, 1303, 1305** with each column including operations that may be executed by a different component of the arrangement **600**. The column **1301** includes operations that can be performed by a user group **614, 616, 618, 620** or, for example, a single user authorized to make changes on behalf of the user group **614, 616, 618, 620**. The column **1303** includes operations executed by the frontend service **612**. The column **1305** includes operations performed by the application orchestration service **604**, such as by the application updater service **603**.

(130) At operation **1302**, the user group **614, 616, 618, 620** begins the upgrade process, for example, by opening and/or otherwise requesting an application dashboard. The application dashboard may be any suitable user interface element provided to the user group **614, 616, 618, 620** by the frontend service **612**. At operation **1304**, the frontend service **612** builds consumer context data for the user group **614, 616, 618, 620**. This may include, for example, querying the user or users from the user group **614, 616, 618, 620** who opened the dashboard at operation **1302** and/or accessing data about the user and/or user group **614, 616, 618, 620**. At operation **1306**, the frontend service **612** requests a filtered list of permissible versions for the user group **614, 616, 618, 620**. The filtered list of permissible versions may indicate versions of the application that are available as target versions for upgrade, as indicated by the upgrade strategy map **610** and permissible for use by the user group **614, 616, 618, 620**, as indicated by the rollout strategy map **609**. The request may be directed to the version catalog service **607**. The version catalog service **607** may store a previously-generated list of permissible applications for the user group **614, 616, 618, 620** and/or may generate or request that the list of permissible versions be generated using the consumer context data generated at operation **1304**.

(131) The frontend service **612** receives the filtered list of permissible versions at operation **1308** and provides the list to the user group at operation **1310**. For example, the frontend service **612** may provide the list of permissible versions via a user interface provided to a user or users of the user group **614, 616, 618, 620**. The user or users from the user group **614, 616, 618, 620** select a version for updating at operation **1312**. In some examples, the upgrade may be a patch, as described herein.

(132) At operation **1314**, the frontend service **612** collects parameters for the selected update and calls the application orchestration service **604** (e.g., the application updater service **603**). The collected parameters may include, for example, an indication of the user group **614, 616, 618, 620** making the request, an indication of the current instance **622, 624, 626, 628** of the application for that user group **614, 616, 618, 620** including, for example, a version of the instance, and an indication of the selected target version.

(133) The application orchestration service **604** (e.g., the application updater service **603**), at operation **1316**, retrieves the container image of the target version referred to at the product version map **608** and, at operation **1318**, upgrades the application instance. Upgrading the application instance may include, for example, spinning up one or more containers to execute a new application instance for the requesting user group **614, 616, 618, 620**. In some examples, upgrading the application instance may also include executing a data migration at a persistence **630, 632, 634, 636** of the requesting user group **614, 616, 618, 620**. For example, the persistence **630, 632, 634, 636** or a portion thereof may be modified from a prior schema to a target schema used by the target version.

(134) At operation **1320**, the frontend service reports to the user group **614, 616, 618, 620** that the

new instance of the application is ready for use. The user group **614, 616, 618, 620** may commence to use the new instance of the application at operation **1322**.

(135) FIG. **14** is a flowchart showing one example of a process flow **1400** that may be executed in the arrangement **600** to maintain the rollout strategy map **609** to make an application version available to a particular user group **614, 616, 618, 620**. The process flow **1400** includes three columns **1401, 1403, 1405** with each column including operations that may be executed by a different component of the arrangement **600**. The column **1401** includes operations that can be performed by an application development team. The application development team may be, for example, a user group **614, 616, 618, 620** that has development and/or administrative responsibilities for an application. The column **1403** includes operations executed by a deployment service, such as the deployment service **638**. The column **1405** includes operations performed by the application orchestration service **604**, such as by the application operator service **605**.

(136) At operation **1402**, the application development team declares a default version for the application. This may include, for example, generating and/or updating one or more rollout records at the rollout strategy map **609**. For example, the deployment service may add a rollout record to the rollout strategy map **609**. The rollout record may not include selector data or may include selector data indicating that the selected default version is a default version that can be provided to different user groups **614, 616, 618, 620**. At operation **1404**, the deployment service updates a rollout strategy map **609** for the application to indicate the default version.

(137) Before the rollout strategy map **609** is updated at operation **1404**, the application orchestration service **604** may not utilize the rollout strategy map **609** to filter versions of the application from the product version map **608**. After the rollout strategy map **609** is updated at operation **1404**, the application orchestration service **604** may begin utilizing the updated rollout strategy map **609**, referred to as V1, at operation **1414**. For example, when the user group **614, 616, 618, 620** provisions a new instance **622, 6234, 626, 628** of the application and/or upgrades a previous instance, the default version may be available.

(138) At operation **1408**, the application development team enables a first version of the application for a user group **614, 616, 618, 620**. The user group **614, 616, 618, 620** may be, for example, an internal user group that is part of a QA team at the entity developing the application. In another example, the user group **614, 616, 618, 620** is a group associated with a customer who is to be authorized to access a beta or other early or pre-release version. At operation **1410**, the application development team enables a version of the application for the user group. The enabled version may be, for example, a beat version or other pre-release version for the user group **614, 616, 618, 620**. In various examples, however, any version of the application may be enabled for the user group. Enabling the version may include updating the rollout strategy map **609** and requesting that the deployment service implement the update at the cloud environment **602**. Before the rollout strategy is enabled at operation **1404**, the product version map may be used, at operation **1406**, without filtering from the rollout strategy map.

(139) At operation **1412**, the deployment service may update the rollout strategy map to another updated rollout strategy map **609** version, called V2 in FIG. **14**. This may include, for example, adding at least one rollout record to the rollout strategy map **609**. The rollout record may include an indication of the version enabled for the user group **614, 616, 618, 620** and selector data indicating consumer context properties of the user group **614, 616, 618, 620**. After the update at operation **1412**, the application orchestration service **604** may begin to use the V2 version of the rollout strategy map **609** at operation **1422**. For example, when the user group **614, 616, 618, 620** requests initiation of an application instance **622, 624, 626, 628** and/or an upgrade to an existing instance, the first version may be available.

(140) At operation **1416**, the application development team may disable the first version of the application for the user group **614, 616, 618, 620**. For example, a bug in the version may have been detected. Also, in some examples, the first version may have been made available for an evaluation

period that has passed. At operation **1418**, the application development team removes the first version of the application for the user group **614, 616, 618, 620**. Disabling the first version may include requesting that the deployment service update the rollout strategy map **609**.

(141) At operation **1420**, the deployment service may update the rollout strategy map to another updated rollout strategy map **609** version, called V3 in FIG. **14**. In some examples, the updated rollout strategy map **609** may omit the rollout record referencing the first version from the rollout strategy map **609** and/or include modified selector data in that rollout record so that it does not refer to the user group **614, 616, 618, 620**. After the update at operation **1420**, the application orchestration service **604** may begin to use the V3 version of the rollout strategy map **609**. For example, when the user group **614, 616, 618, 620** requests initiation of an application instance **622, 624, 626, 628** and/or an upgrade to an existing instance, the first version may no longer be available.

(142) In some examples, updating of the rollout strategy map, for example, as described in FIG. **14**, may be used to expose and then subsequently remove a testing version of the application. For example, updating of the rollout strategy map at operation **1404** may be to add a testing version rollout record to the rollout strategy map along with testing selector data. The testing version rollout record may describe a testing version of the application that is being made available, for example, to a user group **614, 616, 618, 620**. The selector data may indicate the user group or users groups **614, 616, 618, 620** who are to be permitted to use the testing version. After a time, for example, after testing of the test version is complete, the rollout record may be modified again, for example, at operation **1420**, to remove the testing version rollout record from the rollout strategy map.

(143) In various examples, application orchestration service maintain application instances executing at a cloud environment. For example, as described herein, application orchestrations service may manage the execution of one or more containers implementing an application instance. The application orchestration service may start one or more containers to implement an application instances. Also, for example, upon the crash of one or more containers implementing an application instance, application orchestration service may restart or re-launch another version of the container.

(144) In some examples, application orchestration service (directly or indirectly) access version metadata from the version metadata persistence to maintain application instances. For example, when spinning up a new instance of an application, the application orchestration may refer to a version metadata record from the metadata persistence, such as a product version map record, to find an identifier of the container image or images that are used to launch the application.

(145) The application orchestration service may also utilize a link to the container image or images for an application instance while the application instance is running. For example, when one or more of the containers implementing an application instance crashes, the application orchestration service may again utilize an indication of the relevant container image or images to restart the crashed container or containers to maintain the application instance.

(146) In some examples, the application orchestration service may refer back to version metadata persistence, for example, at a product version map, to identify container images used to re-start a crashed container during execution of an application instance. Practically, however, if the application orchestration service uses the version metadata persistence to maintain the continued execution of application instances, it may be necessary to maintain version metadata at the version metadata persistence for each application version (e.g., at the production version map) as long as any instance of that application version is executing at the cloud environment.

(147) The number of actual versions of an application executing in the cloud environment and the number of versions supported for provisioning and upgrading, however may not be the same. For example, it may be desirable to remove version metadata from the version metadata persistence for a particular version of the application to prevent new instances of that version from being provisioned without removing application orchestration support for instances of that particular

version that are still executing at the cloud environment.

(148) Various examples address these and other issues utilizing instance metadata for an application instance. Instance metadata is metadata describing an application instance that is stored at a container management metadata persistence that is associated with the executing instance. In examples where the application orchestration service implements a Kubernetes® arrangement, the instance metadata, and/or addition thereto, may be created using a custom resource.

(149) In various examples, when an application instance is initiated, the application orchestration service writes an indication of one or more container images used for executing the application instance to the instance metadata. The application orchestration service may be further arranged to refer to the instance metadata associated with an application instance to identify container images for maintaining the application instance. In this way, the version metadata relating to an application version may be removed from the version metadata persistence when it is desirable to no longer launch new instances of the version while permitting maintenance of any existing instance of the version that are already executing.

(150) FIG. 15 is a diagram showing one example of an arrangement **1500** for maintaining versions of a legacy on-premises software application in a cloud environment **1502** using a version metadata record. The cloud environment **1502** executes instances of applications, such as a database management application. An example application instance **1522** is shown. Although FIG. 15 shows only a single application instance **1522**, it will be appreciated that the cloud environment **1502** may execute multiple application instances, such as, for example, shown in FIGS. 1 and 6. The cloud environment **1502** also provides various services **1512**, **1504**, **1507** for supporting the application instance **1522** (or instances).

(151) The application instance **1522** may be provided to users from a user group **1514**. The user group **1514** may include a group of one or more users who use the application instance **1522** at the cloud environment **1502**. As described herein, the user group **1514** may be associated with a stakeholder. Although one user group **1514** is shown in FIG. 15, it will be appreciated that some examples may include various different user groups utilizing different application instances, for example, as described herein with respect to FIGS. 1 and 6.

(152) The cloud environment **1502** may include one or more computing devices, such as servers, for executing the application instance **1522** and services **1512**, **1504**, **1507**. In some examples, the cloud environment **1502** is one of a number of cloud environments implemented by a cloud service provider at different geographic locations. For example, the cloud environment **1502** may be implemented at a data center or data centers geographically near to the user group **1514**. Other cloud environments similar to the cloud environment **1502** may be implemented at data centers geographically near to other user groups to provide the application to those user groups.

(153) The cloud environment **1502** also implements a persistence **1530**. For examples, when the application instance **1522** implements a database management system, the persistence **1530** may store data managed by the database management application. The persistence **1530** may be implemented using any suitable data storage device or devices such as, for example, one or more disks, one or more solid state drives, one or more random access memories, etc. In some examples, the software application implemented at the cloud environment **1502** is a database management application for an in-memory database such as, for example, the HANA® and S/4 HANA® systems available from SAP SE of Walldorf, Germany.

(154) The software application provided by the cloud environment **1502** of FIG. 15 may be a legacy on-premises application converted for provision by the cloud environment **1502**. For example, the application instance **1522** may comprise one or more containers including executables, libraries, and/or other components for providing the software application. In some examples, containers for executing the various application instances **1522**, **1524**, **1526**, **1528** are arranged as Kubernetes® clusters, for example, managed and/or maintained by the application orchestration service **1504**.

(155) The cloud environment **1502** also executes one or more services **1507**, **1504**, **1512** for managing execution of the application instance **1522** (and other application instances executing at the cloud environment **1502**). The services **1504**, **1507**, **1512** may execute at the cloud environment **1502** in one or more containers. A frontend service **1512** may interface with the user group **1514**. For example, the frontend service **1512** may provide users from the user group **1514** with a user interface to allow the users to manage the application instance **1522** and/or to launch new application instances. In some examples, the frontend service **1512** provides users from the user group **1514** with data about available target versions to which the user group's executing application instance **1522** can be upgraded, as described herein.

(156) Application orchestration services **1504** may manage the execution of the application instance **152**. The application orchestration service **1504** may be implemented as a single service and/or as a set of one or more subservices. In the example of FIG. 15, the application orchestration services **1504** include an application operator service **1505** and an application updater service **1503**.

(157) The application operator service **1505** may manage the execution of the one or more containers implementing the application instance **1522** (and other application instances executing at the cloud environment **1502**). For example, the application operator service **1505** may begin the execution of one or more containers for implementing the application instance **1522**, for example, from one or more container images associated with the one or more containers. Also, if one or more containers associated with the application instance **1522** crash, the application operator service **1505** may spin up a replacement for the crashed container or containers. The application updater service **1503** may manage application versions.

(158) In the example of FIG. 15, the version metadata persistence **1506** includes a product version map **1508**, an upgrade strategy map **1510** and a rollout strategy map **1509**. The product version map **1508** and upgrade strategy map **1510** may be similar to the product version map **108**, **608** and upgrade strategy map **110**, **610** described herein with respect to FIGS. 1 and 6. The rollout strategy map **1509** may include rollout records, with each rollout record describing a version of the application and selector data, where the selector data describes user group context properties describing user groups that are permitted to use the indicated application version, for example, similar to the rollout strategy map **609** of FIG. 6.

(159) The arrangement **1500** also includes a version catalog service **1507**, which may be similar to the version catalog service **607** of FIG. 6. The version catalog service **1507** may be used by the application orchestration service **1504** and/or the frontend service **1512** to access the version metadata persistence **1506**. For example, the version catalog service **1507** may utilize the product version map **1508**, rollout strategy map **1509**, and/or upgrade strategy map **1510** to generate lists of permissible versions for the user group **1514** and/or for other user groups using the cloud environment **1502**. A list of permissible versions for a user group **1514** indicates versions of the application that are both available and that the user group **1514** is authorized to use. The version catalog service **1507** may generate and/or update a list of permissible versions for the user group **1514** for example, upon a request from the user group **1514** to initiate and/or upgrade the application instance **1522** and/or periodically.

(160) In some examples, the application orchestration service **1504** may request that the version catalog service **1507** utilize the rollout strategy map **1509** in conjunction with the upgrade strategy map **1510**. For example, the list of permissible versions for the user group **1514** at the version catalog service **1507** may include versions that the user group **1514** is permitted to use, as indicated by the rollout strategy map **1509** and that are permissible upgrades from the current version of an application instance **1522** for the user group **1514**. An example rollout strategy map **1509** is provided by EXAMPLE 3.

(161) In the example of FIG. 15, the application instance **1522** comprises instance metadata **1523**. The instance metadata **1523** may be implemented at a container management metadata persistence **1525**. The container management metadata persistence **1525** may be implemented, for example, by

the application orchestration service **1504**. In examples where the application orchestration service **1504** implements a Kubernetes® container orchestration system, the container management metadata persistence **1525** may be implemented using a custom resource. In some examples, the container management metadata persistence **1525** persists if one or more of the containers implementing the application instance **1522** crash. In this way, the application orchestration service **1504** may refer to the instance metadata **1523** to re-start a crashed or otherwise failed container of the application instance **1522**.

(162) The arrangement **1500** of FIG. **15** also includes a deployment service **1538**. The deployment service **1538** may execute at one or more computing devices, such as computing devices implemented by the cloud service provider that implements the cloud environment **1502**. In some examples, the deployment service **1538** is executed at the same computing devices that implement the environment **1502**. The deployment service **1538** may be programmed to launch the cloud environment **1502**, for example, as part of a CI/CD pipeline. In some examples, the deployment service **1538** may also be used to update the metadata persistence **1506**, for example, as described herein. For example, the CI/CD pipeline may be used to generate new versions, patches, etc. for the application instances **1522**, **1524**, **1526**, **1528**. When the deployment service **1538** deploys new version, patch, upgrade, or the like, it may also make appropriate upgrades to the metadata persistence **1506**.

(163) FIG. **16** is a flowchart showing one example of a process flow **1600** that may be executed by in the arrangement **1500** of FIG. **15** to launch the application instance **1522** using instance metadata **1523**. At operation **1602**, the application orchestration service **1504** (e.g., the application operator service **1505** thereof) receives a request to launch the application instance. The request may be received from the user group **1514** and may indicate a version of the application that is to be launched as the requested instance, as described herein.

(164) At operation **1604**, the application orchestration service **1504** accesses a version metadata record for the application. The version metadata record may be, for example, a product version map record stored at the product version map **1508**. The version metadata may include an indication of one or more containers that are to be executed to implement the version of the application. At operation **1606**, the application orchestration service **1504** launches the application instance **1522**. At operation **1608**, the application orchestration service **1504** stores instance metadata **1523** to the container management metadata persistence **1525**.

(165) In this way, the application orchestration service **1504** may utilize the instance metadata **1523** to maintain the application instance **1522** and may not need to refer back to the version metadata during execution of the application instance **1522**.

(166) Consider the example version metadata record given by EXAMPLE 4 below:

Example 4

(167) TABLE-US-00004
apiVersion: products.example.com/v1alpha1 kind: ProductVersionMap
metadata: name: Application 2 spec: enablement: rolloutstrategy: staged-hana-rollout
versions: - name: Application 2 Edition 2019.32 releaseCycle: generally-available-quarterly
track: 2019.32 id: 2019.32.10 version: image-host: <docker repository> image-name:
<application image name> image-tag: 2019.32.10 build-id: 4.00.000.32.10 release-notes:
<url to release notes> - name: Application Edition 2019.36 releaseCycle: evaluation track:
2019.36 id: 2019.36.7 version: image-host: <docker repository> image-name: <application
image name> image-tag: 2019.36.7 build-id: 4.00.000.36.7 release-notes: <url to release
notes> expiration-date: "2029-11-30"

The version metadata record illustrated by EXAMPLE 4 may be a product version map record or other suitable record from the version metadata persistence **1506**. The version metadata record of EXAMPLE 4 indicates two different versions of an application called Application 2 having different maturity levels. A maturity level field called "generally-available-quarterly" indicates a version of the application called "4.00.000.32.10." A maturity level field called "evaluation"

indicates a version of the application called “4.00.000.36.7.” Each maturity level field includes version metadata for the referred-to versions including, for example, an image host field indicating a location of one or more container images for implementing the version, and an image name field indicating the name of the one or more container images for implementing the version. In EXAMPLE 4, the version metadata for a version also indicates an image tag, a build identifier, and a link to release notes.

(168) When the user group **1514** requests an instance **1522** of the application, for example, as described at operation **1602**, the application orchestration service **1504** may copy the version metadata for the requested application version. The copied version metadata may be stored to the instance metadata **1523**, for example, as described at operation **1608**. EXAMPLE 5 below shows example instance metadata for the application instance **1522**:

Example 5

(169) TABLE-US-00005

apiVersion:	application.example.com/v1alpha1	kind:	Application
metadata:	name: application2-instance	spec:	product-version: releaseCycle: generally-available-quarterly track: 2019.32 . . . status: <u>current-product-version:</u> <u>name:</u>
	<u>Application Edition 2019.32</u>		<u>releaseCycle: generally-available-quarterly</u> <u>track: 2019.32</u>
	<u>id: 2019.32.10</u>	<u>version:</u>	<u>image-host: <docker repository></u> <u>image-name:</u>
	<u><application image name></u>	<u>image-tag: 2019.32.10</u>	<u>build-id: 4.00.000.32.10</u>
	<u>release-notes: <url to release notes></u>	. . .	

In EXAMPLE 5, the application orchestration service **1504** has copied version metadata, indicated by the underlined portion, to the instance metadata **1523** for the application instance **1522**. The version metadata includes, for example, an indication of at least one container used to implement the version of the application (e.g., image-name, image-tag).

(170) FIG. 17 is a flowchart showing one example of a process flow **1700** that may be executed in the arrangement **1500** of FIG. 15 to update the version metadata persistence **1506**. The process flow **1700** is executed by one or more services **1504**, **1512**, **1507**, **1538** of the arrangement **1500**. In some examples, the process flow **1700** is executed by the deployment service **1538** upon instructions from a user group, such as the user group **1514** and/or an administrative user group.

(171) At operation **1702**, the at least one service receives a new version of the application. Receiving the new version of the application may include, for example, receiving one or more container images that may be executed to implement an instance of the new version. In some examples, receiving the new version also includes receiving version metadata for the new version. The version metadata may include, for example, an indication of the at least one container image and a maturity level of new version.

(172) At operation **1704**, the at least one service adds to the version metadata persistence **1506** a version metadata record for the new application version. This may include, for example, identifying the version metadata corresponding to the indicated maturity level and writing the new version data to that location. At operation **1706**, the at least one service removes version metadata related to the previous version (e.g., the previous version having the same maturity level). In this way, the version metadata corresponding to the new version replaces the corresponding version metadata corresponding to the previous version (e.g., the previous version having the same maturity level). Removal of the version metadata for the previous version may, as described herein, remove the previous version for the versions available to the user group **1514** for launch or upgrade.

(173) Because the version metadata for the previous version may have been copied to the instance metadata for already-executing instances of the previous version, the previous versions may continue to be supported by the application orchestration service **1504** after their corresponding version metadata is removed from the version metadata persistence **1506**. For example, FIG. 18 is a process flow **1800** that may be executed in the arrangement **1500** of FIG. 15, e.g., by the application orchestration services, to maintain an application instance. At operation **1804**, the application orchestration service **1504** (or other suitable cloud service) determines if one or more of

the containers implementing the application instance **1522** is down or has crashed. If not, the application orchestration service **1504** may continue to check whether a container of the application instance **1522** has crashed. For example, the application orchestration service **1504** may check the application instance **1522** at operation **1804** periodically and/or in response to an alert or other indication of that a container may have crashed.

(174) If one or more of the containers implementing the application instance **1522** has crashed, the application orchestration service **1504**, at operation **1802**, accesses version metadata for the version of the application instance **1522** from the instance metadata **1523** at the container management metadata persistence **1525**. The instance metadata **1523**, as described herein, may include an indication of one or more container images for restarting the crashed container. The application orchestration service **1504** may use the indicated container image or images to relaunch the crashed container at operation **1806**. In this way, the application orchestration service **1504** may maintain the application instance **1522** without referring to the version metadata persistence **1506** after the launch of the application instance **1522**.

(175) FIG. **19** is a flowchart showing one example of a process flow **1900** for updating the application instance **1522** in the example arrangement **1500** of FIG. **15**. The process flow **1900**, in some examples, is executed by the application orchestration service **1504**, e.g., the application updater service **1503** thereof.

(176) At operation **1902**, the application orchestration service **1504** receives a request to upgrade the application instance **1522**. The request may be received in any suitable manner. In some examples, the user group **1514** requests the upgrade, as described herein. In other examples, an administrative user group may use the deploy service **1538** to modify a version metadata record associated with the application instance **1522** to indicate a new version of the application. For example, the product version map record or other suitable version metadata record indicating the same maturity level as the application instance **1522** may be updated to refer a new version of the application (e.g., a version different than the version executing at the application instance **1522**).

(177) Consider again the example product version map record given by EXAMPLE 4 above. In EXAMPLE 4, the maturity level field corresponding to “generally-available quarterly” indicates an application version “4.00.000.32.10.” An example of the product version map record of EXAMPLE 4 after the execution of the operation **1902** is given by EXAMPLE 6 below:

Example 6

(178) TABLE-US-00006 apiVersion: products.example.com/v1alpha1 kind: ProductVersionMap
metadata: name: Application 2 spec: enablement: rolloutstrategy: staged-hana-rollout
versions: - name: Application 2 Edition 2019.34 releaseCycle: generally-available-quarterly
track: 2019.34 id: 2019.34.1 version: image-host: <docker repository> image-name:
<application image name> image-tag: 2019.34.1 build-id: 4.00.000.34.1 release-notes: <url
to release notes> - name: Application Edition 2019.36 releaseCycle: evaluation track: 2019.36
id: 2019.36.7 version: image-host: <docker repository> image-name: <application image
name> image-tag: 2019.36.7 build-id: 4.00.000.36.7 release-notes: <url to release notes>
expiration-date: “2029-11-30”

The underlined portion of EXAMPLE 6 indicates the maturity level field corresponding to “generally-available-quarterly” reflecting an update from the reference to the version “4.00.000.32.10” at EXAMPLE 4 to the reference to a version “4.00.000.34.1” at EXAMPLE 6. In some examples, this modification to the product version map record or other metadata record may prompt an upgrade of the application instance **1522** from the previous version to the new version (e.g., in EXAMPLES 4 and 6 from version “4.00.000.32.10” to version “4.00.000.34.1”).

(179) At operation **1906**, the application orchestration service **1504** may upgrade the application instance **1522**, for example, as described herein. At operation **1906**, the application orchestration service **1504** may store new instance metadata indicating at least one container for implementing the new version of the application. In this way, the instance metadata referring to the previous

version of the application instance **1522** may be deleted.

(180) Consider again EXAMPLES 4, 5, and 6 herein. Recall that EXAMPLE 5 shows the instance metadata **1523** for the application instance **1522**. EXAMPLE 7 below shows an example of the instance metadata for the application instance **1522** after an upgrade to a new version (e.g., from version “4.00.000.32.10” indicated by the version metadata record of EXAMPLE 4 to version “4.00.000.34.1” indicated by the version metadata record of EXAMPLE 6).

Example 7

(181) TABLE-US-00007 apiVersion: application.example.com/v1alpha1 kind: Application
metadata: name: application2-instance spec: product-version: releaseCycle: generally-
available-quarterly track: 2019.32 ... status: current-product-version: name:
Application 2 Edition 2019.34 releaseCycle: generally-available-quarterly track:
2019.34 id: 2019.34.1 version: image-host: <docker repository>
image-name: <application image name> image-tag: 2019.34.1 build-id:
4.00.000.34.1 release-notes: <url to release notes> ...

(182) In EXAMPLE 7, the underlined portion indicates a change relative to EXAMPLE 5 and shows a change from version “4.00.000.32.10” indicated by the version metadata record of EXAMPLE 4 to version “4.00.000.34.1” indicated by the version metadata record of EXAMPLE 6.

(183) In view of the disclosure above, various examples are set forth below. It should be noted that one or more features of an example, taken in isolation or combination, should be considered within the disclosure of this application.

Examples

(184) Example 1 is a method for upgrading a cloud-implemented database management application, comprising: receiving, by a first cloud service executing in a cloud environment, a request from a first user group to access the database management application; accessing, by the first cloud service, consumer context data comprising a plurality of context properties of the first user group; accessing, by the first cloud service, a rollout strategy map describing a set of rollout records, the set of rollout records comprising a first rollout record describing a first version of the database management application and first selector data; comparing, by the first cloud service, the plurality of context properties of the first user group to the first selector data; based on the comparing, adding the first version of the database management application to a list of permissible versions for the first user group; providing, by the first cloud service, the list of permissible versions for the first user group to at least one user associated with the first user group; accessing, by the first cloud service, an indication that the first user group has selected the first version; and initiating execution of a first instance of the first version for the first user group at the cloud environment, the first instance to access a first data persistence associated with the first user group.

(185) In Example 2, the subject matter of Example 1 optionally includes the initiating execution of the first instance comprising: accessing a container image associated with the first version of the database management application; and using the container image to initiate a container at the cloud environment, the container comprising at least one executable for providing the first version of the database management application.

(186) In Example 3, the subject matter of any one or more of Examples 1-2 optionally includes accessing, by the first cloud service, a version map for the database management application, the version map indicating a list of available versions of the database management application, the version map referencing the rollout strategy map.

(187) In Example 4, the subject matter of any one or more of Examples 1-3 optionally includes the request being a request to upgrade the database management application for the first user group, the method further comprising: determining, by the first cloud service, a first version of the database management application used by the first user group; accessing, by the first cloud service, an upgrade strategy map; determining, using the upgrade strategy map, a set of target versions of the database management application suitable for upgrading from the first version; determining that a

second version of the database management application is not part of the set of target versions; and removing the second version from the list of permissible versions for the first user group before providing the list of permissible versions for the first user group to the at least one user associated with the first user group.

(188) In Example 5, the subject matter of any one or more of Examples 1-4 optionally includes accessing, by the first cloud service, an indication of a default version of the database management application; and modifying the rollout strategy map to include a default rollout record, the default rollout record describing the default version of the database management application.

(189) In Example 6, the subject matter of any one or more of Examples 1-5 optionally includes accessing, by the first cloud service, an instruction to enable a testing version of the database management application for the first user group; and modifying the rollout strategy map to include a testing version rollout record, the testing version rollout record describing the testing version of the database management application and testing selector data.

(190) In Example 7, the subject matter of Example 6 optionally includes comparing, by the first cloud service, the plurality of context properties of the first user group to the testing selector data; and based on comparing, adding the testing version of the database management application to the list of permissible versions for the first user group.

(191) In Example 8, the subject matter of any one or more of Examples 6-7 optionally includes accessing, by the first cloud service, an instruction to disable the testing version of the database management application for the first user group; and modifying the rollout strategy map to disable the test version for the first user group.

(192) In Example 9, the subject matter of Example 8 optionally includes the modifying of the rollout strategy map to disable the test version for the first user group comprises at least one of modifying the testing selector data and deleting the test version rollout record.

(193) Example 10 is a system for upgrading a cloud-implemented database management application, comprising: a cloud environment comprising at least one computing device, the at least one computing device programmed to perform operations comprising: receiving, by a first cloud service executing in a cloud environment, a request from a first user group to access the database management application; accessing, by the first cloud service, consumer context data comprising a plurality of context properties of the first user group; accessing, by the first cloud service, a rollout strategy map describing a set of rollout records, the set of rollout records comprising a first rollout record describing a first version of the database management application and first selector data; comparing, by the first cloud service, the plurality of context properties of the first user group to the first selector data; based on the comparing, adding the first version of the database management application to a list of permissible versions for the first user group; providing, by the first cloud service, the list of permissible versions for the first user group to at least one user associated with the first user group; accessing, by the first cloud service, an indication that the first user group has selected the first version; and initiating execution of a first instance of the first version for the first user group at the cloud environment, the first instance to access a first data persistence associated with the first user group.

(194) In Example 11, the subject matter of Example 10 optionally includes the initiating execution of the first instance comprising: accessing a container image associated with the first version of the database management application; and using the container image to initiate a container at the cloud environment, the container comprising at least one executable for providing the first version of the database management application.

(195) In Example 12, the subject matter of any one or more of Examples 10-11 optionally includes the operations further comprising accessing, by the first cloud service, a version map for the database management application, the version map indicating a list of available versions of the database management application, the version map referencing the rollout strategy map.

(196) In Example 13, the subject matter of any one or more of Examples 10-12 optionally includes

the request being a request to upgrade the database management application for the first user group, the operations further comprising: determining, by the first cloud service, a first version of the database management application used by the first user group; accessing, by the first cloud service, an upgrade strategy map; determining, using the upgrade strategy map, a set of target versions of the database management application suitable for upgrading from the first version; determining that a second version of the database management application is not part of the set of target versions; and removing the second version from the list of permissible versions for the first user group before providing the list of permissible versions for the first user group to the at least one user associated with the first user group.

(197) In Example 14, the subject matter of any one or more of Examples 10-13 optionally includes the operations further comprising: accessing, by the first cloud service, an indication of a default version of the database management application; and modifying the rollout strategy map to include a default rollout record, the default rollout record describing the default version of the database management application.

(198) In Example 15, the subject matter of any one or more of Examples 10-14 optionally includes the operations further comprising: accessing, by the first cloud service, an instruction to enable a testing version of the database management application for the first user group; and modifying the rollout strategy map to include a testing version rollout record, the testing version rollout record describing the testing version of the database management application and testing selector data.

(199) In Example 16, the subject matter of Example 15 optionally includes the operations further comprising: comparing, by the first cloud service, the plurality of context properties of the first user group to the testing selector data; and based on comparing, adding the testing version of the database management application to the list of permissible versions for the first user group.

(200) In Example 17, the subject matter of any one or more of Examples 15-16 optionally includes the operations further comprising: accessing, by the first cloud service, an instruction to disable the testing version of the database management application for the first user group; and modifying the rollout strategy map to disable the test version for the first user group.

(201) In Example 18, the subject matter of Example 17 optionally includes the modifying of the rollout strategy map to disable the test version for the first user group comprises at least one of modifying the testing selector data and deleting the test version rollout record.

(202) Example 19 is a non-transitory machine-readable medium comprising instructions thereon that, when executed by at least one processor at a cloud environment, causes the at least one processor to perform operations comprising: receiving, by a first cloud service executing in a cloud environment, a request from a first user group to access a database management application; accessing, by the first cloud service, consumer context data comprising a plurality of context properties of the first user group; accessing, by the first cloud service, a rollout strategy map describing a set of rollout records, the set of rollout records comprising a first rollout record describing a first version of the database management application and first selector data; comparing, by the first cloud service, the plurality of context properties of the first user group to the first selector data; based on the comparing, adding the first version of the database management application to a list of permissible versions for the first user group; providing, by the first cloud service, the list of permissible versions for the first user group to at least one user associated with the first user group; accessing, by the first cloud service, an indication that the first user group has selected the first version; and initiating execution of a first instance of the first version for the first user group at the cloud environment, the first instance to access a first data persistence associated with the first user group.

(203) In Example 20, the subject matter of Example 19 optionally includes the initiating execution of the first instance comprising: accessing a container image associated with the first version of the database management application; and using the container image to initiate a container at the cloud environment, the container comprising at least one executable for providing the first version of the

database management application.

(204) FIG. **20** is a block diagram **2000** showing one example of a software architecture **2002** for a computing device. The architecture **2002** may be used in conjunction with various hardware architectures, for example, as described herein. FIG. **20** is merely a non-limiting example of a software architecture and many other architectures may be implemented to facilitate the functionality described herein. A representative hardware layer **2004** is illustrated and can represent, for example, any of the above referenced computing devices. In some examples, the hardware layer **2004** may be implemented according to the architecture of the computer system of FIG. **20**.

(205) The representative hardware layer **2004** comprises one or more processing units **2006** having associated executable instructions **2008**. Executable instructions **2008** represent the executable instructions of the software architecture **2002**, including implementation of the methods, modules, subsystems, and components, and so forth described herein and may also include memory and/or storage modules **2010**, which also have executable instructions **2008**. Hardware layer **2004** may also comprise other hardware as indicated by other hardware **2012** which represents any other hardware of the hardware layer **2004**, such as the other hardware illustrated as part of the architecture **2002**.

(206) In the example architecture of FIG. **20**, the software architecture **2002** may be conceptualized as a stack of layers where each layer provides particular functionality. For example, the software architecture **2002** may include layers such as an operating system **2014**, libraries **2016**, frameworks/middleware **2018**, applications **2020**, and presentation layer **2044**. Operationally, the applications **2020** and/or other components within the layers may invoke API calls **2024** through the software stack and access a response, returned values, and so forth illustrated as messages **2026** in response to the API calls **2024**. The layers illustrated are representative in nature and not all software architectures have all layers. For example, some mobile or special purpose operating systems may not provide a frameworks/middleware layer **2018**, while others may provide such a layer. Other software architectures may include additional or different layers.

(207) The operating system **2014** may manage hardware resources and provide common services. The operating system **2014** may include, for example, a kernel **2028**, services **2030**, and drivers **2032**. The kernel **2028** may act as an abstraction layer between the hardware and the other software layers. For example, the kernel **2028** may be responsible for memory management, processor management (e.g., scheduling), component management, networking, security settings, and so on. The services **2030** may provide other common services for the other software layers. In some examples, the services **2030** include an interrupt service. The interrupt service may detect the receipt of an interrupt and, in response, cause the architecture **2002** to pause its current processing and execute an interrupt service routine (ISR) when an interrupt is accessed.

(208) The drivers **2032** may be responsible for controlling or interfacing with the underlying hardware. For instance, the drivers **2032** may include display drivers, camera drivers, Bluetooth® drivers, flash memory drivers, serial communication drivers (e.g., Universal Serial Bus (USB) drivers), Wi-Fi® drivers, NFC drivers, audio drivers, power management drivers, and so forth depending on the hardware configuration.

(209) The libraries **2016** may provide a common infrastructure that may be utilized by the applications **2020** and/or other components and/or layers. The libraries **2016** typically provide functionality that allows other software modules to perform tasks in an easier fashion than to interface directly with the underlying operating system **2014** functionality (e.g., kernel **2028**, services **2030** and/or drivers **2032**). The libraries **2016** may include system **2034** libraries (e.g., C standard library) that may provide functions such as memory allocation functions, string manipulation functions, mathematic functions, and the like. In addition, the libraries **2016** may include API libraries **2036** such as media libraries (e.g., libraries to support presentation and manipulation of various media format such as MPEG4, H.264, MP3, AAC, AMR, JPG, PNG),

graphics libraries (e.g., an OpenGL framework that may be used to render 2D and 3D in a graphic content on a display), database libraries (e.g., SQLite that may provide various relational database functions), web libraries (e.g., WebKit that may provide web browsing functionality), and the like. The libraries **2016** may also include a wide variety of other libraries **2038** to provide many other APIs to the applications **2020** and other software components/modules.

(210) The frameworks **2018** (also sometimes referred to as middleware) may provide a higher-level common infrastructure that may be utilized by the applications **2020** and/or other software components/modules. For example, the frameworks **2018** may provide various graphic user interface (GUI) functions, high-level resource management, high-level location services, and so forth. The frameworks **2018** may provide a broad spectrum of other APIs that may be utilized by the applications **2020** and/or other software components/modules, some of which may be specific to a particular operating system or platform.

(211) The applications **2020** includes built-in applications **2040** and/or third-party applications **2042**. Examples of representative built-in applications **2040** may include, but are not limited to, a contacts application, a browser application, a book reader application, a location application, a media application, a messaging application, and/or a game application. Third-party applications **2042** may include any of the built-in applications **2040** as well as a broad assortment of other applications. In a specific example, the third-party application **2042** (e.g., an application developed using the Android™ or iOS™ software development kit (SDK) by an entity other than the vendor of the particular platform) may be mobile software running on a mobile operating system such as iOS™, Android™, Windows® Phone, or other mobile computing device operating systems. In this example, the third-party application **2042** may invoke the API calls **2024** provided by the mobile operating system such as operating system **2014** to facilitate functionality described herein.

(212) The applications **2020** may utilize built-in operating system functions (e.g., kernel **2028**, services **2030** and/or drivers **2032**), libraries (e.g., system **2034**, APIs **2036**, and other libraries **2038**), and frameworks/middleware **2018** to create user interfaces to interact with users of the system. Alternatively, or additionally, in some systems interactions with a user may occur through a presentation layer, such as presentation layer **2044**. In these systems, the application/module “logic” can be separated from the aspects of the application/module that interact with a user.

(213) Some software architectures utilize virtual machines. In the example of FIG. **20**, this is illustrated by virtual machine **2048**. A virtual machine creates a software environment where applications/modules can execute as if they were executing on a hardware computing device. A virtual machine is hosted by a host operating system (operating system **2014**) and typically, although not always, has a virtual machine monitor **2046**, which manages the operation of the virtual machine as well as the interface with the host operating system (i.e., operating system **2014**). A software architecture executes within the virtual machine such as an operating system **2050**, libraries **2052**, frameworks/middleware **2054**, applications **2056** and/or presentation layer **2058**. These layers of software architecture executing within the virtual machine **2048** can be the same as corresponding layers previously described or may be different.

(214) Modules, Components and Logic

(215) Certain embodiments are described herein as including logic or a number of components, modules, or mechanisms. Modules may constitute either software modules (e.g., code embodied (1) on a non-transitory machine-readable medium or (2) in a transmission signal) or hardware-implemented modules. A hardware-implemented module is a tangible unit capable of performing certain operations and may be configured or arranged in a certain manner. In example embodiments, one or more computer systems (e.g., a standalone, client, or server computer system) or one or more hardware processors may be configured by software (e.g., an application or application portion) as a hardware-implemented module that operates to perform certain operations as described herein.

(216) In various embodiments, a hardware-implemented module may be implemented

mechanically or electronically. For example, a hardware-implemented module may comprise dedicated circuitry or logic that is permanently configured (e.g., as a special-purpose processor, such as a field programmable gate array (FPGA) or an application-specific integrated circuit (ASIC)) to perform certain operations. A hardware-implemented module may also comprise programmable logic or circuitry (e.g., as encompassed within a general-purpose processor or another programmable processor) that is temporarily configured by software to perform certain operations. It will be appreciated that the decision to implement a hardware-implemented module mechanically, in dedicated and permanently configured circuitry, or in temporarily configured circuitry (e.g., configured by software) may be driven by cost and time considerations.

(217) Accordingly, the term “hardware-implemented module” should be understood to encompass a tangible entity, be that an entity that is physically constructed, permanently configured (e.g., hardwired), or temporarily or transitorily configured (e.g., programmed) to operate in a certain manner and/or to perform certain operations described herein. Considering embodiments in which hardware-implemented modules are temporarily configured (e.g., programmed), each of the hardware-implemented modules need not be configured or instantiated at any one instance in time. For example, where the hardware-implemented modules comprise a general-purpose processor configured using software, the general-purpose processor may be configured as respective different hardware-implemented modules at different times. Software may accordingly configure a processor, for example, to constitute a particular hardware-implemented module at one instance of time and to constitute a different hardware-implemented module at a different instance of time.

(218) Hardware-implemented modules can provide information to, and receive information from, other hardware-implemented modules. Accordingly, the described hardware-implemented modules may be regarded as being communicatively coupled. Where multiple of such hardware-implemented modules exist contemporaneously, communications may be achieved through signal transmission (e.g., over appropriate circuits and buses that connect the hardware-implemented modules). In embodiments in which multiple hardware-implemented modules are configured or instantiated at different times, communications between such hardware-implemented modules may be achieved, for example, through the storage and retrieval of information in memory structures to which the multiple hardware-implemented modules have access. For example, one hardware-implemented module may perform an operation, and store the output of that operation in a memory device to which it is communicatively coupled. A further hardware-implemented module may then, at a later time, access the memory device to retrieve and process the stored output. Hardware-implemented modules may also initiate communications with input or output devices, and can operate on a resource (e.g., a collection of information).

(219) The various operations of example methods described herein may be performed, at least partially, by one or more processors that are temporarily configured (e.g., by software) or permanently configured to perform the relevant operations. Whether temporarily or permanently configured, such processors may constitute processor-implemented modules that operate to perform one or more operations or functions. The modules referred to herein may, in some example embodiments, comprise processor-implemented modules.

(220) Similarly, the methods described herein may be at least partially processor-implemented. For example, at least some of the operations of a method may be performed by one or more processors or processor-implemented modules. The performance of certain of the operations may be distributed among the one or more processors, not only residing within a single machine, but deployed across a number of machines. In some example embodiments, the processor or processors may be located in a single location (e.g., within a home environment, an office environment, or a server farm), while in other embodiments the processors may be distributed across a number of locations.

(221) The one or more processors may also operate to support performance of the relevant operations in a “cloud computing” environment or as a “software as a service” (SaaS). For

example, at least some of the operations may be performed by a group of computers (as examples of machines including processors), these operations being accessible via a network (e.g., the Internet) and via one or more appropriate interfaces (e.g., APIs).

(222) Electronic Apparatus and System

(223) Example embodiments may be implemented in digital electronic circuitry, or in computer hardware, firmware, or software, or in combinations of them. Example embodiments may be implemented using a computer program product, e.g., a computer program tangibly embodied in an information carrier, e.g., in a machine-readable medium for execution by, or to control the operation of, data processing apparatus, e.g., a programmable processor, a computer, or multiple computers.

(224) A computer program can be written in any form of programming language, including compiled or interpreted languages, and it can be deployed in any form, including as a standalone program or as a module, subroutine, or other unit suitable for use in a computing environment. A computer program can be deployed to be executed on one computer or on multiple computers at one site or distributed across multiple sites and interconnected by a communication network.

(225) In example embodiments, operations may be performed by one or more programmable processors executing a computer program to perform functions by operating on input data and generating output. Method operations can also be performed by, and apparatus of example embodiments may be implemented as, special purpose logic circuitry, e.g., an FPGA or an ASIC.

(226) The computing system can include clients and servers. A client and server are generally remote from each other and typically interact through a communication network. The relationship of client and server arises by virtue of computer programs running on the respective computers and having a client-server relationship to each other. In embodiments deploying a programmable computing system, it will be appreciated that both hardware and software architectures merit consideration. Specifically, it will be appreciated that the choice of whether to implement certain functionality in permanently configured hardware (e.g., an ASIC), in temporarily configured hardware (e.g., a combination of software and a programmable processor), or in a combination of permanently and temporarily configured hardware may be a design choice. Below are set out hardware (e.g., machine) and software architectures that may be deployed, in various example embodiments.

(227) Example Machine Architecture and Machine-Readable Medium

(228) FIG. **21** is a block diagram of a machine in the example form of a computer system **2100** within which instructions **2124** may be executed for causing the machine to perform any one or more of the methodologies discussed herein. In alternative embodiments, the machine operates as a standalone device or may be connected (e.g., networked) to other machines. In a networked deployment, the machine may operate in the capacity of a server or a client machine in server-client network environment, or as a peer machine in a peer-to-peer (or distributed) network environment. The machine may be a personal computer (PC), a tablet PC, a set-top box (STB), a personal digital assistant (PDA), a cellular telephone, a web appliance, a network router, switch, or bridge, or any machine capable of executing instructions (sequential or otherwise) that specify actions to be taken by that machine. Further, while only a single machine is illustrated, the term “machine” shall also be taken to include any collection of machines that individually or jointly execute a set (or multiple sets) of instructions to perform any one or more of the methodologies discussed herein.

(229) The example computer system **2100** includes a processor **2102** (e.g., a central processing unit (CPU), a graphics processing unit (GPU), or both), a main memory **2104**, and a static memory **2106**, which communicate with each other via a bus **2108**. The computer system **2100** may further include a video display unit **2110** (e.g., a liquid crystal display (LCD) or a cathode ray tube (CRT)). The computer system **2100** also includes an alphanumeric input device **2112** (e.g., a keyboard or a touch-sensitive display screen), a user interface (UI) navigation (or cursor control) device **2114** (e.g., a mouse), a disk drive unit **2116**, a signal generation device **2118** (e.g., a speaker), and a

network interface device **2120**.

(230) Machine-Readable Medium

(231) The disk drive unit **2116** includes a machine-readable medium **2122** on which is stored one or more sets of data structures and instructions **2124** (e.g., software) embodying or utilized by any one or more of the methodologies or functions described herein. The instructions **2124** may also reside, completely or at least partially, within the main memory **2104** and/or within the processor **2102** during execution thereof by the computer system **2100**, with the main memory **2104** and the processor **2102** also constituting machine-readable media **2122**.

(232) While the machine-readable medium **2122** is shown in an example embodiment to be a single medium, the term “machine-readable medium” may include a single medium or multiple media (e.g., a centralized or distributed database, and/or associated caches and servers) that store the one or more instructions **2124** or data structures. The term “machine-readable medium” shall also be taken to include any tangible medium that is capable of storing, encoding, or carrying instructions **2124** for execution by the machine and that cause the machine to perform any one or more of the methodologies of the present disclosure, or that is capable of storing, encoding, or carrying data structures utilized by or associated with such instructions **2124**. The term “machine-readable medium” shall accordingly be taken to include, but not be limited to, solid-state memories, and optical and magnetic media. Specific examples of machine-readable media **2122** include non-volatile memory, including by way of example semiconductor memory devices, e.g., erasable programmable read-only memory (EPROM), electrically erasable programmable read-only memory (EEPROM), and flash memory devices; magnetic disks such as internal hard disks and removable disks; magneto-optical disks; and CD-ROM and DVD-ROM disks.

(233) Transmission Medium

(234) The instructions **2124** may further be transmitted or received over a communications network **2126** using a transmission medium. The instructions **2124** may be transmitted using the network interface device **2120** and any one of a number of well-known transfer protocols (e.g., HTTP). Examples of communication networks include a local area network (LAN), a wide area network (WAN), the Internet, mobile telephone networks, plain old telephone (POTS) networks, and wireless data networks (e.g., WiFi and WiMax networks). The term “transmission medium” shall be taken to include any intangible medium that is capable of storing, encoding, or carrying instructions **2124** for execution by the machine, and includes digital or analog communications signals or other intangible media to facilitate communication of such software.

(235) Although an embodiment has been described with reference to specific example embodiments, it will be evident that various modifications and changes may be made to these embodiments without departing from the broader spirit and scope of the disclosure. Accordingly, the specification and drawings are to be regarded in an illustrative rather than a restrictive sense. The accompanying drawings that form a part hereof show by way of illustration, and not of limitation, specific embodiments in which the subject matter may be practiced. The embodiments illustrated are described in sufficient detail to enable those skilled in the art to practice the teachings disclosed herein. Other embodiments may be utilized and derived therefrom, such that structural and logical substitutions and changes may be made without departing from the scope of this disclosure. This Detailed Description, therefore, is not to be taken in a limiting sense, and the scope of various embodiments is defined only by the appended claims, along with the full range of equivalents to which such claims are entitled.

(236) Such embodiments of the inventive subject matter may be referred to herein, individually and/or collectively, by the term “invention” merely for convenience and without intending to voluntarily limit the scope of this application to any single invention or inventive concept if more than one is in fact disclosed. Thus, although specific embodiments have been illustrated and described herein, it should be appreciated that any arrangement calculated to achieve the same purpose may be substituted for the specific embodiments shown. This disclosure is intended to

cover any and all adaptations or variations of various embodiments. Combinations of the above embodiments, and other embodiments not specifically described herein, will be apparent to those of skill in the art upon reviewing the above description.

Claims

1. A method for upgrading a database management application, comprising: receiving, by a service executing in a cloud environment, a request from a user group to access the database management application; accessing, by the service, data describing a plurality of properties of the user group; accessing, by the service, map data describing a set of map data records, the set of map data records comprising a record describing a first version of the database management application and selector data; comparing, by the service, the plurality of properties of the user group to the selector data; based on the comparing, adding the first version of the database management application to a list of versions for the user group; providing, by the service, the list of versions for the user group to at least one user associated with the user group; accessing, by the service, an indication that the user group has selected the first version of the database management application; and initiating execution of an instance of the first version of the database management application for the user group at the cloud environment, the instance to access a data persistence associated with the user group.
2. The method of claim 1, the initiating execution of the instance comprising: accessing a container image associated with the first version of the database management application; and using the container image to initiate a container at the cloud environment, the container comprising at least one executable for providing the first version of the database management application.
3. The method of claim 1, further comprising accessing, by the service, a version map for the database management application, the version map indicating a list of available versions of the database management application, the version map referencing the map data.
4. The method of claim 1, the request being a request to upgrade the database management application for the user group, the method further comprising: determining, by the service, the first version of the database management application used by the user group; accessing, by the service, an upgrade strategy map; determining, using the upgrade strategy map, a set of target versions of the database management application suitable for upgrading from the first version of the database management application; determining that a second version of the database management application is not part of the set of target versions; and removing the second version from the list of versions for the user group before providing the list of versions for the user group to the at least one user associated with the user group.
5. The method of claim 1, further comprising: accessing, by the service, an indication of a default version of the database management application; and modifying the map data to include a default record, the default record describing the default version of the database management application.
6. The method of claim 1, further comprising: accessing, by the service, an instruction to enable a testing version of the database management application for the user group; and modifying the map data to include a testing version record, the testing version record describing the testing version of the database management application and testing selector data.
7. The method of claim 6, further comprising: comparing, by the service, the plurality of properties of the user group to the testing selector data; and based on comparing, adding the testing version of the database management application to the list of versions for the user group.
8. The method of claim 6, further comprising: accessing, by the service, an instruction to disable the testing version of the database management application for the user group; and modifying the map data to disable the testing version for the user group.
9. The method of claim 8, the modifying of the map data to disable the testing version for the user group comprises at least one of modifying the testing selector data and deleting the testing version

record.

10. A system for upgrading a database management application, comprising: a cloud environment comprising at least one computing device, the at least one computing device programmed to perform operations comprising: receiving, by a service executing in a cloud environment, a request from a user group to access the database management application; accessing, by the service, data describing a plurality of properties of the user group; accessing, by the service, a map data describing a set of map data records, the set of map data records comprising a record describing a first version of the database management application and selector data; comparing, by the service, the plurality of properties of the user group to the selector data; based on the comparing, adding the first version of the database management application to a list of versions for the user group; providing, by the service, the list of versions for the user group to at least one user associated with the user group; accessing, by the service, an indication that the user group has selected the first version of the database management application; and initiating execution of an instance of the first version of the database management application for the user group at the cloud environment, the instance to access a data persistence associated with the user group.

11. The system of claim 10, the initiating execution of the instance comprising: accessing a container image associated with the first version of the database management application; and using the container image to initiate a container at the cloud environment, the container comprising at least one executable for providing the first version of the database management application.

12. The system of claim 10, the operations further comprising accessing, by the service, a version map for the database management application, the version map indicating a list of available versions of the database management application, the version map referencing the map data.

13. The system of claim 10, the request being a request to upgrade the database management application for the user group, the operations further comprising: determining, by the service, the first version of the database management application used by the user group; accessing, by the service, an upgrade strategy map; determining, using the upgrade strategy map, a set of target versions of the database management application suitable for upgrading from the first version; determining that a second version of the database management application is not part of the set of target versions; and removing the second version from the list of versions for the user group before providing the list of versions for the user group to the at least one user associated with the user group.

14. The system of claim 10, the operations further comprising: accessing, by the service, an indication of a default version of the database management application; and modifying the map data to include a default record, the default record describing the default version of the database management application.

15. The system of claim 10, the operations further comprising: accessing, by the service, an instruction to enable a testing version of the database management application for the user group; and modifying the map data to include a testing version record, the testing version record describing the testing version of the database management application and testing selector data.

16. The system of claim 15, the operations further comprising: comparing, by the service, the plurality of properties of the user group to the testing selector data; and based on comparing, adding the testing version of the database management application to the list of versions for the user group.

17. The system of claim 15, the operations further comprising: accessing, by the service, an instruction to disable the testing version of the database management application for the user group; and modifying the map data to disable the testing version for the user group.

18. The system of claim 17, the modifying of the map data to disable the testing version for the user group comprises at least one of modifying the testing selector data and deleting the testing version record.

19. A non-transitory machine-readable medium comprising instructions thereon that, when

executed by at least one processor at a cloud environment, causes the at least one processor to perform operations comprising: receiving, by a service executing in a cloud environment, a request from a user group to access a database management application; accessing, by the service, consumer context data comprising a plurality of properties of the user group; accessing, by the service, a map data describing a set of map data records, the set of map data records comprising a record describing a first version of the database management application and selector data; comparing, by the service, the plurality of properties of the user group to the selector data; based on the comparing, adding the first version of the database management application to a list of versions for the user group; providing, by the service, the list of versions for the user group to at least one user associated with the user group; accessing, by the service, an indication that the user group has selected the first version of the database management application; and initiating execution of an instance of the first version of the database management application for the user group at the cloud environment, the instance to access a data persistence associated with the user group.

20. The non-transitory machine-readable medium of claim 19, the initiating execution of the instance comprising: accessing a container image associated with the first version of the database management application; and sing the container image to initiate a container at the cloud environment, the container comprising at least one executable for providing the first version of the database management application.
