

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent
Kind Code
Date of Patent
Inventor(s)

12399718
B2
August 26, 2025
Shemy; Regev et al.

Apparatuses, methods, and systems for hashing instructions

Abstract

Systems, methods, and apparatuses relating to performing hashing operations on packed data elements are described. In one embodiment, a processor includes a decode circuit to decode a single instruction into a decoded single instruction, the single instruction including at least one first field that identifies eight 32-bit state elements A, B, C, D, E, F, G, and H for a round according to a SM3 hashing standard and at least one second field that identifies an input message; and an execution circuit to execute the decoded single instruction to: rotate state element C left by 9 bits to form a rotated state element C, rotate state element D left by 9 bits to form a rotated state element D, rotate state element G left by 19 bits to form a rotated state element G, rotate state element H left by 19 bits to form a rotated state element H, perform two rounds according to the SM3 hashing standard on the input message and state element A, state element B, rotated state element C, rotated state element D, state element E, state element F, rotated state element G, and rotated state element H to generate an updated state element A, an updated state element B, an updated state element E, and an updated state element F, and store the updated state element A, the updated state element B, the updated state element E, and the updated state element F into a location specified by the single instruction.

Inventors: Shemy; Regev (Kiryat Ata, IL), Sperber; Zeev (Zichron Yackov, IL), Feghali; Wajdi (Boston, MA), Gopal; Vinodh (Westborough, MA), Gradstein; Amit (Binyamina, IL), Rubanovich; Simon (Haifa, IL), Gulley; Sean (Sudbury, MA), Albrekht; Ilya (Tempe, AZ), Doweck; Jacob (Haifa, IL), Yallouz; Jose (Haifa, IL), Anati; Ittai (Ramat Hasharon, IL)

Applicant: Intel Corporation (Santa Clara, CA)

Family ID: 1000008780642

Assignee: Intel Corporation (Santa Clara, CA)

Appl. No.: 18/336985

Filed: June 17, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240036865 A1	Feb. 01, 2024

Related U.S. Application Data

continuation parent-doc US 17688728 20220307 US 11681530 20230620 child-doc US 18336985
continuation parent-doc US 17537373 20211129 US 11567772 20230131 child-doc US 17688728
continuation parent-doc US 17087536 20201102 US 11188335 20211130 child-doc US 17537373
continuation parent-doc US 16370459 20190329 US 10824428 20201103 child-doc US 17087536

Publication Classification

Int. Cl.: G06F9/30 (20180101); G06F9/38 (20180101); H04L9/06 (20060101)

U.S. Cl.:

CPC G06F9/30145 (20130101); G06F9/30036 (20130101); G06F9/30038 (20230801);
G06F9/30043 (20130101); G06F9/30196 (20130101); G06F9/3887 (20130101);
G06F9/3888 (20230801); H04L9/0643 (20130101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7684563	12/2009	Olson et al.	N/A	N/A
8160242	12/2011	Grinchuk et al.	N/A	N/A
8856546	12/2013	Gueron et al.	N/A	N/A
9027104	12/2014	Wolrich et al.	N/A	N/A
9658854	12/2016	Wolrich	N/A	G06F 9/30036
9985648	12/2017	Wu et al.	N/A	N/A
10148428	12/2017	Wolrich	N/A	G06F 9/30145
10152326	12/2017	Yap et al.	N/A	N/A
10503510	12/2018	Wolrich et al.	N/A	N/A
10686591	12/2019	Wolrich et al.	N/A	N/A
2005/0089160	12/2004	Crispin et al.	N/A	N/A
2010/0250966	12/2009	Olson	712/205	G06F 9/3836
2011/0142228	12/2010	Crispin et al.	N/A	N/A
2012/0128149	12/2011	Boersma et al.	N/A	N/A
2014/0093069	12/2013	Wolrich et al.	N/A	N/A
2014/0195782	12/2013	Yap et al.	N/A	N/A
2015/0006856	12/2014	Rash et al.	N/A	N/A
2015/0186139	12/2014	Wolrich	712/208	H04L 9/0643

2015/0280917	12/2014	Wolrich et al.	N/A	N/A
2016/0070931	12/2015	Gueron	380/28	G06F 9/30007
2016/0092688	12/2015	Wolrich	713/190	G06F 9/30036
2017/0255469	12/2016	Wolrich et al.	N/A	N/A
2019/0095202	12/2018	Malladi et al.	N/A	N/A
2022/0188114	12/2021	Shemy et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
101894229	12/2009	CN	N/A
107094369	12/2016	CN	N/A
108959123	12/2017	CN	N/A
109508206	12/2018	CN	N/A
114510272	12/2021	CN	N/A
114546490	12/2021	CN	N/A
3189617	12/2016	EP	N/A
3343357	12/2017	EP	N/A
2016/036491	12/2015	WO	N/A

OTHER PUBLICATIONS

ARM Architecture Reference Manual, “ARMv8 and ARM v8-A architecture profile,” ARM DDI 0487D.a, 110 Fulbourn Road, Cambridge, England CB1 9NJ, retrieved from [https://static.docs.arm.com/ddi0487/da/DDI0487D_a_armv8_arm.pdf ?](https://static.docs.arm.com/ddi0487/da/DDI0487D_a_armv8_arm.pdf?_ga=2.118486921.1618531128.155259426...)

European Search Report and Search Opinion, EP App. No. 20156066.1, Aug. 25, 2020, 9 pages. cited by applicant

European Search Report and Search Opinion, EP App. No. 22156135.0, Mar. 28, 2022, 11 pages. cited by applicant

European Search Report and Search Opinion, EP App. No. 22161283.1, Mar. 28, 2022, 10 pages. cited by applicant

FIPS PUB 180-4, “Secure Hash Standard (SHS),” Information Technology Laboratory, National Institute of Standards and Technology, publication available from <http://dx.doi.org/10.6028/NIST.FIPS.180-4>, Aug. 2015, 36 pages. cited by applicant

Intel, “Intel 64 and IA-32 Architectures Software Developer's Manual”, Combined vols. 1, 2A, 2B, 2C, 2D, 3A, 3B, 3C, 3D and 4, Order No. 325462-069US, <https://software.intel.com/sites/default/files/managed/39/c5/325462-sdm-vol-1-2abcd-3abcd.pdf>, Jan. 2019, 4898 pages. cited by applicant

Notice of Allowance, U.S. Appl. No. 16/370,459, Jun. 25, 2020, 10 pages. cited by applicant

Notice of Allowance, U.S. Appl. No. 17/087,536, Aug. 5, 2021, 10 pages. cited by applicant

Notice of Allowance, U.S. Appl. No. 17/537,373, Nov. 14, 2022, 2 pages. cited by applicant

Notice of Allowance, U.S. Appl. No. 17/537,373, Oct. 6, 2022, 10 pages. cited by applicant

Notice of Allowance, U.S. Appl. No. 17/688,728, Feb. 13, 2023, 10 pages. cited by applicant

Notice of Allowance, U.S. Appl. No. 17/688,728, May 17, 2023, 2 pages. cited by applicant

Office Action, EP App. No. 20156066.1, Apr. 1, 2022, 5 pages. cited by applicant

Office Action, EP App. No. 20156066.1, Aug. 5, 2022, 4 pages. cited by applicant

Office Action, EP App. No. 22156135.0, Feb. 20, 2023, 4 pages. cited by applicant

Office Action, EP App. No. 22161283.1, Feb. 20, 2023, 4 pages. cited by applicant

Shen S., et al., “The SM3 Cryptographic Hash Function draft-sca-cfrg-sm3-02,” Network Working Group, Internet-Draft, <https://tools.ietf.org/html/draft-sca-cfrg-sm3-02>, Jan. 8, 2018, 71 pages. cited by applicant

Wikipedia, "SHA-2," Retrieved from <https://en.wikipedia.org/wiki/SHA-2> on Mar. 26, 2019, 8 pages. cited by applicant
Intention to Grant, EP App. No. 22156135.0, Jun. 26, 2024, 6 pages. cited by applicant
Intention to Grant, EP App. No. 22161283.1, Jun. 19, 2024, 6 pages. cited by applicant
Intention to Grant, EP App. No. 20156066.1, Feb. 5, 2024, 6 pages. cited by applicant
Bin et al., "Implementation of Extended Instruction Set for AES Fast Algorithm", vol. 40 No. 6, Jun. 2012, 7 pages. cited by applicant
First Office Action, CN App. No. 202210136333.1, Apr. 18, 2025, 13 pages (6 pages of English Translation and 7 pages of Original Document). cited by applicant
Li et al., "Efficient Implementation for Hash Function SHA-512", vol. 33 No.7, Apr. 2007, 4 pages. cited by applicant
Office Action, CN App. No. 202210347963.3, Jun. 5, 2025, 07 pages of Original Document only. cited by applicant

Primary Examiner: Sun; Michael

Attorney, Agent or Firm: NICHOLSON DE VOS WEBSTER & ELLIOTT LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) The present application is a continuation application claiming priority from U.S. patent application Ser. No. 17/688,728 filed Mar. 7, 2022, now U.S. Pat. No. 11,681,530, which is a continuation application claiming priority from U.S. patent application Ser. No. 17/537,373 filed Nov. 29, 2021, now U.S. Pat. No. 11,567,7724, which is a continuation application claiming priority from U.S. patent application Ser. No. 17/087,536 filed Nov. 2, 2020, now U.S. Pat. No. 11,188,335, which is a continuation application claiming priority from U.S. patent application Ser. No. 16/370,459 filed Mar. 29, 2019, now U.S. Pat. No. 10,824,428, each of which is incorporated herein by reference in its entirety.

TECHNICAL FIELD

(1) The disclosure relates generally to electronics, and, more specifically, an embodiment of the disclosure relates to processor circuitry for hashing operations.

BACKGROUND

(2) A processor, or set of processors, executes instructions from an instruction set, e.g., the instruction set architecture (ISA). The instruction set is the part of the computer architecture related to programming, and generally includes the native data types, instructions, register architecture, addressing modes, memory architecture, interrupt and exception handling, and external input and output (I/O). It should be noted that the term instruction herein may refer to a macro-instruction, e.g., an instruction that is provided to the processor for execution, or to a micro-instruction, e.g., an instruction that results from a processor's decode circuit decoding macro-instructions.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The present disclosure is illustrated by way of example and not limitation in the figures of the accompanying drawings, in which like references indicate similar elements and in which:
(2) FIG. 1 illustrates a hardware processor coupled to a memory according to embodiments of the disclosure.

- (3) FIG. 2 illustrates a hardware processor coupled to storage that includes one or more hashing instructions according to embodiments of the disclosure.
- (4) FIG. 3 illustrates a compression function of a SM3 hashing standard according to embodiments of the disclosure.
- (5) FIG. 4A illustrates a section of pseudocode to perform hashing according to a SM3 hashing standard according to embodiments of the disclosure.
- (6) FIG. 4B illustrates another section of pseudocode to perform hashing according to a SM3 hashing standard according to embodiments of the disclosure.
- (7) FIG. 5 illustrates a method of processing a hashing instruction according to a SM3 hashing standard according to embodiments of the disclosure.
- (8) FIG. 6 illustrates circuitry including an execution circuit with rotation circuitry and SM3 hashing round circuitry according to embodiments of the disclosure.
- (9) FIG. 7 illustrates a method of processing a pre-rotation instruction according to a SM3 hashing standard according to embodiments of the disclosure.
- (10) FIG. 8 illustrates circuitry including an execution circuit with rotation circuitry according to embodiments of the disclosure.
- (11) FIG. 9 illustrates a method of processing a post-rotation instruction according to a SM3 hashing standard according to embodiments of the disclosure.
- (12) FIG. 10 illustrates circuitry including an execution circuit with rotation circuitry according to embodiments of the disclosure.
- (13) FIG. 11 illustrates a method of processing an intermediate message instruction according to a SM3 hashing standard according to embodiments of the disclosure.
- (14) FIG. 12 illustrates circuitry including an execution circuit with intermediate SM3 message calculation circuitry according to embodiments of the disclosure.
- (15) FIG. 13 illustrates intermediate SM3 message calculation circuitry according to embodiments of the disclosure.
- (16) FIG. 14 illustrates a method of processing a final message instruction according to a SM3 hashing standard according to embodiments of the disclosure.
- (17) FIG. 15 illustrates circuitry including an execution circuit with final SM3 message calculation circuitry according to embodiments of the disclosure.
- (18) FIG. 16 illustrates final SM3 message calculation circuitry according to embodiments of the disclosure.
- (19) FIG. 17 illustrates SHA512 hashing circuitry according to embodiments of the disclosure.
- (20) FIG. 18 illustrates a method of processing a hashing instruction according to a SHA512 hashing standard according to embodiments of the disclosure.
- (21) FIG. 19 illustrates circuitry including an execution circuit with SHA512 hashing round circuitry according to embodiments of the disclosure.
- (22) FIG. 20 illustrates a method of processing an intermediate message instruction according to a SHA512 hashing standard according to embodiments of the disclosure.
- (23) FIG. 21 illustrates circuitry including an execution circuit with intermediate SHA512 message calculation circuitry according to embodiments of the disclosure.
- (24) FIG. 22 illustrates a method of processing a final message instruction according to a SHA512 hashing standard according to embodiments of the disclosure.
- (25) FIG. 23 illustrates circuitry including an execution circuit with final SHA512 message calculation circuitry according to embodiments of the disclosure.
- (26) FIG. 24A is a block diagram illustrating a generic vector friendly instruction format and class A instruction templates thereof according to embodiments of the disclosure.
- (27) FIG. 24B is a block diagram illustrating the generic vector friendly instruction format and class B instruction templates thereof according to embodiments of the disclosure.
- (28) FIG. 25A is a block diagram illustrating fields for the generic vector friendly instruction

formats in FIGS. 24A and 24B according to embodiments of the disclosure.

(29) FIG. 25B is a block diagram illustrating the fields of the specific vector friendly instruction format in FIG. 25A that make up a full opcode field according to one embodiment of the disclosure.

(30) FIG. 25C is a block diagram illustrating the fields of the specific vector friendly instruction format in FIG. 25A that make up a register index field according to one embodiment of the disclosure.

(31) FIG. 25D is a block diagram illustrating the fields of the specific vector friendly instruction format in FIG. 25A that make up the augmentation operation field **2450** according to one embodiment of the disclosure.

(32) FIG. **26** is a block diagram of a register architecture according to one embodiment of the disclosure

(33) FIG. 27A is a block diagram illustrating both an exemplary in-order pipeline and an exemplary register renaming, out-of-order issue/execution pipeline according to embodiments of the disclosure.

(34) FIG. 27B is a block diagram illustrating both an exemplary embodiment of an in-order architecture core and an exemplary register renaming, out-of-order issue/execution architecture core to be included in a processor according to embodiments of the disclosure.

(35) FIG. 28A is a block diagram of a single processor core, along with its connection to the on-die interconnect network and with its local subset of the Level 2 (L2) cache, according to embodiments of the disclosure.

(36) FIG. 28B is an expanded view of part of the processor core in FIG. 28A according to embodiments of the disclosure.

(37) FIG. 29 is a block diagram of a processor that may have more than one core, may have an integrated memory controller, and may have integrated graphics according to embodiments of the disclosure.

(38) FIG. 30 is a block diagram of a system in accordance with one embodiment of the present disclosure.

(39) FIG. 31 is a block diagram of a more specific exemplary system in accordance with an embodiment of the present disclosure.

(40) FIG. 32 is a block diagram of a second more specific exemplary system in accordance with an embodiment of the present disclosure.

(41) FIG. 33 is a block diagram of a system on a chip (SoC) in accordance with an embodiment of the present disclosure.

(42) FIG. 34 is a block diagram contrasting the use of a software instruction converter to convert binary instructions in a source instruction set to binary instructions in a target instruction set according to embodiments of the disclosure.

DETAILED DESCRIPTION

(43) In the following description, numerous specific details are set forth. However, it is understood that embodiments of the disclosure may be practiced without these specific details. In other instances, well-known circuits, structures and techniques have not been shown in detail in order not to obscure the understanding of this description.

(44) References in the specification to “one embodiment,” “an embodiment,” “an example embodiment,” etc., indicate that the embodiment described may include a particular feature, structure, or characteristic, but every embodiment may not necessarily include the particular feature, structure, or characteristic. Moreover, such phrases are not necessarily referring to the same embodiment. Further, when a particular feature, structure, or characteristic is described in connection with an embodiment, it is submitted that it is within the knowledge of one skilled in the art to affect such feature, structure, or characteristic in connection with other embodiments whether or not explicitly described.

(45) A (e.g., hardware) processor (e.g., having one or more cores) may execute instructions (e.g., a thread of instructions) to operate on data, for example, to perform a hashing operation or operations. For example, software may request a hashing operation and a hardware processor (e.g., a core or cores thereof) may perform the operation in response to the request. One non-limiting example of a hashing operation is a round (or rounds) of hashing according to a (e.g., cryptographic) hashing standard. A hashing standard may include one or more hash functions that take a message as an input, generate a corresponding hash value or digest by applying the hash function(s) to the message, and output the hash value or digest. The hash value may include an updated state word. Typically, the same hash value (or digest) should be generated if the same hash function is evaluated with the same message. Such hash functions are used for various purposes, such as for verification (e.g., verifying the integrity of files, data, or messages), identification (e.g., identifying files, data, or messages), authentication (e.g., generating message authentication codes), generating digital signatures, generating pseudorandom numbers, etc. As one illustrative example, a hash function is used to generate a hash value for a given (input) message. At a later time, a hash value may be recomputed for the given message using the same hash function. If the hash values are identical, then it can be assumed that the message has not been changed. In contrast, if the hash values are different, then it can be assumed that the message has been changed.

(46) One example of a hashing standard is a SM3 hashing standard. A SM3 hashing standard has been published by the Chinese Commercial Cryptography Association Office and approved by the Chinese government. An English language description of a SM3 hashing standard has been published as the Internet Engineering Task Force (IETF) and titled “The SM3 Cryptographic Hash Function” by S. Shen et al, on Jan. 8, 2018. In one embodiment, a SM3 hashing standard takes an input message as 512-bit blocks (e.g., with 32-bit elements) and generates a 256-bit hash (or digest) value. Certain embodiments herein provide for an instruction or instructions (e.g., as discussed in reference to FIGS. 3-16) that are improvements to the functioning of a processor (e.g., of a computer) itself. For example, including dedicated instruction(s) for a SM3 hashing standard reduces power consumption, reduces the total number of instructions to be decoded/executed to generate a hash (or digest) output value, and increases the performance (e.g., speed of execution) of a hashing operation according to a SM3 hashing standard. Instruction decode circuitry (e.g., a decoder) not having such an instruction as a part of its instruction set would not decode as discussed herein. An execution circuit not having such an instruction as a part of its instruction set would not execute as discussed herein.

(47) Another example of a hashing standard is a SHA512 (or SHA384) hashing standard. A SHA512 (and SHA384) hashing standard has been published in August 2015 by the National Institute of Standards and Technology (NIST) as “Federal Information Processing Standard (FIPS) 180-4 and titled “Secure Hash Standard (SHS)”. In one embodiment, a SHA512 hashing standard takes an input message as 1024-bit blocks (e.g., with 64-bit elements) and generates a 512-bit hash (or message digest) value. Certain embodiments herein provide for an instruction or instructions (e.g., as discussed in reference to FIGS. 17-23) that are improvements to the functioning of a processor (e.g., of a computer) itself. For example, including dedicated instruction(s) for a SHA512 hashing standard reduces power consumption, reduces the total number of instructions to be decoded/executed to generate a hash (or digest) output value, and increases the performance (e.g., speed of execution) of a hashing operation according to a SHA512 hashing standard. Instruction decode circuitry (e.g., a decoder) not having such an instruction as a part of its instruction set would not decode as discussed herein. An execution circuit not having such an instruction as a part of its instruction set would not execute as discussed herein.

(48) FIG. 1 illustrates a hardware processor **100** coupled to a memory **110** according to embodiments of the disclosure. Depicted hardware processor **100** includes a hardware decode circuit **102** (e.g., decode unit or decoder) and a hardware execution circuit **104** (e.g., execution unit). Depicted hardware processor **100** includes register(s) **106**. Registers may include one or more

registers to perform operations in, e.g., additionally or alternatively to access of (e.g., load or store) data in memory **110**. Note that the figures herein may not depict all data communication connections. One of ordinary skill in the art will appreciate that this is to not obscure certain details in the figures. Note that a double headed arrow in the figures may not require two-way communication, for example, it may indicate one-way communication (e.g., to or from that component or device). Any or all combinations of communications paths may be utilized in certain embodiments herein.

(49) Hardware decode circuit **102** may receive an (e.g., single) instruction (e.g., macro-instruction) and decode the instruction, e.g., into micro-instructions and/or micro-operations. Hardware execution circuit **104** may execute the decoded instruction (e.g., macro-instruction) to perform an operation or operations. For example, an instruction to be decoded by decode circuit **102** and for the decoded instruction to be executed by execution circuit **104** may be any instruction discussed herein, e.g., in FIGS. **3-23**. Hardware execution circuit **104** may be any of the execution circuits in the Figures. In certain embodiment, hardware execution circuit **104** is (e.g., only) a hashing execution circuit, for example, that is utilized separately from a general execution circuit or execution unit (e.g., one that performs non-hashing specific instructions).

(50) The decode circuit **102**, execution circuit **104**, and registers **106** may be of a single core of the processor, e.g., and multiple cores each with an instance of the circuitry may be included. The processor (e.g., and core thereof) may be a processor and/or core according to any of the disclosure herein.

(51) FIG. **2** illustrates a hardware processor **200** coupled to storage **202** that includes one or more hashing instructions **204** according to embodiments of the disclosure. In certain embodiments, a hashing instruction is according to any of the disclosure herein. In one embodiment, e.g., in response to a request to perform an operation, the instruction (e.g., macro-instruction) is fetched from storage **202** and sent to decode circuit **206**. In the depicted embodiment, the decode circuit **206** (e.g., decoder) decodes the instruction into a decoded instruction (e.g., one or more micro-instructions or micro-operations). The decoded instruction is then sent for execution, e.g., via scheduler circuit **208** to schedule the decoded instruction for execution.

(52) In certain embodiments, (e.g., where the processor/core supports out-of-order (OoO) execution), the processor includes a register rename/allocator circuit coupled to register file/memory circuit **210** (e.g., unit) to allocate resources and perform register renaming on registers (e.g., vector registers associated with a logical operation and test instruction). In certain embodiments, (e.g., for out-of-order execution), the processor includes one or more scheduler circuits **208** coupled to the decode circuit. The scheduler circuit(s) may schedule one or more operations associated with decoded instructions, including one or more operations decoded from a logical operation and test instruction, for execution on the execution circuit **212**.

(53) In certain embodiments, a write back circuit **214** is included to write back results of an instruction to a destination (e.g., write them to a register(s) and/or memory), for example, so those results are visible within a processor (e.g., visible outside of the execution circuit that produced those results).

(54) One or more of these components (e.g., decode circuit **206**, register rename/register allocator/scheduler **208**, execution circuit **212**, register file/memory **210**, or write back circuit **214**) may be in a single core of a hardware processor (e.g., and multiple cores each with an instance of these components). Hashing instruction **204** may be dedicated instruction(s) for a SM3 hashing standard. Hashing instruction **204** may be dedicated instruction(s) for a SHA512 hashing standard.

(55) Instructions for a SM3 Hashing Standard

(56) In certain embodiments, an SM3 hashing function accepts a message as input. The message may represent a bit string of arbitrary length. The SM3 hashing function performs a number of operations using the input message and generates a hash value or digest having a length of 256-bits after padding and iterative compression.

(57) FIG. 3 illustrates a compression function 312 of a SM3 hashing standard according to embodiments of the disclosure.

(58) Initially, the 256-bit state value $V(i)$ is partitioned into eight 32-bit state elements (e.g., words) A, B, C, D, E, F, G, and H. The initial state value $V(0)$ for the first iteration is a constant defined by the SM3 hashing function. The state words A through H are specified in big-endian format according to the SM3 hashing function but their format in an implementation may vary if desired.

(59) An iterative procedure is then performed on the sequence of blocks. The depicted SM3 hashing standard includes sixty-four iterations or “rounds” (i.e., from j ranging from 0 to 63). As shown, a single round 313 includes a number of different operations. The leftward pointing arrow symbol ($\leftarrow$) represents storing, assigning, or equating the value or parameter on the right to the value or parameter on the left. The symbol “<<<” represents a left rotate operation followed by a number of bits to rotate. The symbol “>>>” represents a right rotate operation followed by a number of bits to rotate. The symbol of the encircled plus sign ($\oplus$) represents a logical exclusive OR (XOR) operation. T_j is a constant having a value as specified in the SM3 hash function that depends on the iteration (i.e., the value of j). For example, T_j may have value. The variables SS1, SS2, TT1, and TT2 are internal 79cc4519 for $0 \leq j \leq 15$ and the value 7a879d8a for $16 \leq j \leq 63$ intermediate values used in the iterations.

(60) FF_j is a Boolean function which varies with round number (j) according to Function 1:

(61) $FF_j(X, Y, Z) = X \text{ XOR } Y \text{ XOR } Z$ ($0 \leq j \leq 15$); or Function1
 $= (X \text{ AND } Y) \text{ OR } (X \text{ AND } Z) \text{ OR } (Y \text{ AND } Z)$ ($16 \leq j \leq 63$)

(62) GG_j is a Boolean function which varies with round number (j) according to Function 2:

(63) $GG_j(X, Y, Z) = X \text{ XOR } Y \text{ XOR } Z$ ($0 \leq j \leq 15$); or Function2
 $= (X \text{ AND } Y) \text{ OR } (\text{NOT } X \text{ AND } Z)$ ($16 \leq j \leq 63$)

(64) P_0 is a permutation function in compression function according to Function 3:

$P_0(X) = X \text{ XOR } (X \ll 9) \text{ XOR } (X \ll 17)$ Function 3

(65) Notice that the term W_j is added to the evaluation of the Boolean function GG_j . Also, the term W'_j is added to the evaluation of the Boolean function FF_j . The terms W_j and W'_j represent message terms, message inputs, or simply messages. For iterations 0 to 15, the terms W_0 to W_{15} are obtained from the 512-bit block being compressed. In particular, the 512-bit message block being compressed is divided or partitioned into sixteen 32-bit words referenced in big-endian format as W_0 to W_{15} . The remaining messages W_j and W'_j are calculated during a message extension or message expansion portion of the SM3 hashing function.

(66) The W_j messages for iterations 16 to 67 may be calculated according to the following Function 4:

$W_j = P_1(W_{j-16} \text{ XOR } W_{j-9} \text{ XOR } (W_{j-3} \ll 15)) \text{ XOR } (W_{j-13} \ll 7) \text{ XOR } W_{j-6}$ Function 4

(67) In Function 4, $P_1(X)$ is a permutation function for message expansion that is defined by the following Function 5:

$P_1(X) = X \text{ XOR } (X \ll 15) \text{ XOR } (X \ll 23)$ Function 5

(68) The W_j messages for iterations 16 to 67 may be calculated according to Function 4 with the permutation function P_1 according to Function 5. Notice that the calculation of a W_j message for a given round (e.g., round j) depend on messages from earlier rounds. In particular, as can be readily seen in Function 4, the W_j message for a given round (e.g., round j) depends on the prior round messages W_{j-16} , W_{j-13} , W_{j-9} , W_{j-6} , and W_{j-3} . W_{j-3} is the message from three rounds back relative to round j , W_{j-16} is the message from sixteen rounds back relative to round j , and so on.

(69) The W'_j messages may be calculated or derived from the W_j messages according to the following Function 6:

$W'_j = W_j \text{ XOR } W_{j+4}$ Function 6

(70) Notice that the W'_j message depends on the W_j message from the same round as well as on the W_{j+4} message from four rounds ahead. Since W_0 to W_{15} are divided or obtained from the 512-bit

message block being compressed, messages W'0 to W'11 may be determined using Function 6 based on the initially known messages W0 to W15. The remaining messages W'12 to W'63 may be determined from messages W16 to W67, which may be calculated using Functions 4-5. Notice that W64 to W67 may be calculated, even though they are not input directly into the compression function, but are needed to calculate W'60 to W'63.

(71) As shorthand herein, a value may include a round indicator with it (e.g., W[0] for round 0 of a value of a word or W[63] for round sixty-three of a value of a word) or not, but it should be understood that hashing by rounds allows for those values to be referenced by a round. Thus, in this disclosure, a value indicator may be assigned to a particular round. As one example, values W[i], W[i+1], W[i+4], W[i+5] where index i indicates the current round. For example, a first round (i=0) including the values (e.g., data words) of W[0], W[1], W[4] and W[5], and a third round (i=3) including the values W[2], W[3], W[6], and W[7] (e.g., each time 4 different message words). As another example, values A-H may be particular values for each round (e.g., "i"), for example, A[i], B[i], C[i], D[i], E[i], G[i], H[i] or A.sub.i, B.sub.i, C.sub.i, D.sub.i, E.sub.i, F.sub.i, G.sub.i, H.sub.i depending on the naming format of the round references).

(72) FIG. 4A illustrates a section of pseudocode **400** to perform hashing according to a SM3 hashing standard according to embodiments of the disclosure. Pseudocode **400** is to perform hashing operations according to a SM3 hashing standard for each data block (e.g., which may be 512-bits, i.e., 64 bytes). Thus, for each block **402**, a message calculation **404** (e.g., one or more message calculation instructions) is performed, followed by two rounds of hashing operations **406**, a rotation of state words C, D, G, and H **408** (e.g., for that particular round), then additional sections of two rounds of hashing operations **410** and rotation of state words C, D, G, and H **412** (e.g., for that particular round), followed by a logical XOR (e.g., for each bit position) according to **414** and **416**. Logical XOR operation produces an output of true (e.g., logical one) for each section (e.g., bit position) if an odd number of its multiple arguments (e.g., in a same section of each of the values being compared) are true, and false (e.g., logical zero) otherwise. In one embodiment, each rotation (**408**, **412**, etc.) performed requires its own instruction, and thus consumes the resources of a set of stages of an instruction pipeline, e.g., fetch, decode, schedule, execution, and commit. In certain embodiments, the rotations (e.g., **408**, **412**) are repeated (e.g., sixteen times) until all (e.g., 64) rounds are complete per data block. For example, according to a SM3 hashing standard, each input block may include 64 bytes of input and thus will be using 32 calls for an instruction that performs two rounds of hashing operations (e.g., **406** or **410**) with a corresponding 32 calls for a rotation instruction (e.g., **408** or **412**).

(73) In one embodiment, block **402** calculation is a block digest calculation that includes 64 rounds and message word calculations and rounds performed. In one embodiment, message calculation (e.g., **404**) calculates a plurality (e.g., four) message words using either the input word or previous calculated words, calculating the next plurality (e.g., four) words relevant for the next rounds calculations. In certain embodiments, there are two (e.g., parallel) workloads: (i) message word calculation (e.g., using VSM3MSG1 and VSM3MSG2 instructions or input loading) for calculating the message words, and (ii) rounds calculation that uses the calculated message words. In one embodiment, these operations are iterated through for each round (e.g., other than the first 16 words read by input, and not by performing the VSM3MSG1 and VSM3MSG2 instructions).

(74) FIG. 4B illustrates another section of pseudocode **401** to perform hashing according to a SM3 hashing standard according to embodiments of the disclosure. As compared to the pseudocode **400** in FIG. 4A, the pseudocode **401** in FIG. 4B includes a single instruction that performs both rotation of state elements C, D, G, and H and two rounds of hashing operations. Thus, for each block **405**, a message calculation **407** (e.g., from the input data block or one or more message calculation instructions as discussed herein) is performed, followed by two rounds of hashing operations **409** that also perform (e.g., before the hashing operations) a rotation of state words C, D, G, and H, then one or more instances of additional instructions that each perform two rounds of hashing operations

and rotation (e.g., before the hashing operations) of state words C, D, G, and H (**411**), followed by a logical XOR (e.g., for each bit position) according to **413** and **415** (e.g., for each single block of data). In one embodiment, the old value (e.g., OLD_ABEF and OLD_CDGH) is sourced from the first value of A-H (e.g., when starting each block such at the beginning the first A-H values are the initial values defined by A SM3 hashing standard).

(75) For example, according to a SM3 hashing standard, each input block may include 64 bytes of input and thus will be using 32 calls for an instruction that performs two rounds of hashing operations that also performs (e.g., before the hashing operations) a rotation of state words C, D, G, and H (e.g., **409** or **411**) without utilizing the processor resources for a corresponding 32 calls for a rotation instruction (e.g., **408** or **412** in FIG. 4A). Thus, embodiments of an instruction that performs two rounds of SM3 hashing operations that also performs (e.g., before the hashing operations) a rotation of SM3 state words C, D, G, and H reduces the total number of instructions in pseudocode **401** which saves power, processing resources, and time by removing the decoding, scheduling, committing, etc. for the removed 32 calls for a rotation instruction (e.g., **408** or **412** in FIG. 4A), e.g., for each block with additional pre rotation **403** and post rotation **417** called when iterating through the loop. Examples of an instruction that performs two rounds of SM3 hashing operations that also performs (e.g., before the hashing operations) a rotation of SM3 state words C, D, G, and H are discussed below in reference to FIGS. 5 and 6. Pre-rotate instruction **403** is discussed below in reference to FIGS. 7-8 and post-rotate instruction **417** is discussed below in reference to FIGS. 9-10.

(76) In one embodiment, block **405** calculation is a block digest calculation that includes 64 rounds and message word calculations and rounds performed. In one embodiment, message calculation (e.g., **407**) calculates a plurality (e.g., four) message words using either the input word or previous calculated words, calculating the next plurality (e.g., four) words relevant for the next rounds calculations. In certain embodiments, there are two (e.g., parallel) workloads: (i) message word calculation (e.g., using VSM3MSG1 and VSM3MSG2 instructions or input loading) for calculating the message words, and (ii) rounds calculation that uses the calculated message words. In one embodiment, these operations are iterated through for each round (e.g., other than the first 16 words read by input, and not by performing the VSM3MSG1 and VSM3MSG2 instructions).

(77) FIG. 5 illustrates a method **500** of processing a (e.g., two rounds of) hashing instruction according to a SM3 hashing standard according to embodiments of the disclosure. A processor (e.g., or processor core) may perform method **500**, e.g., in response to receiving a request to execute an instruction from software. Depicted method **500** includes processing a SM3 hashing standard instruction by: fetch a single instruction having an opcode that indicates two rounds of hashing according to a SM3 hashing standard are to be performed and at least one first field that identifies eight (e.g., 32-bit) state elements A, B, C, D, E, F, G, and H for a round (e.g., “i”) according to the SM3 hashing standard and at least one second field that identifies an input message (e.g., and a location to write a resultant) at **502**, decode the instruction into a decoded instruction at **504**, retrieve data associated with the identified source operands at **506**, (optionally) schedule the decoded instruction for execution at **508**, execute the decoded instruction to rotate state element C left by 9 bits to form a rotated state element C, rotate state element D left by 9 bits to form a rotated state element D, rotate state element G left by 19 bits to form a rotated state element G, rotate state element H left by 19 bits to form a rotated state element H, and (e.g., after those rotations) perform two rounds according to the SM3 hashing standard on the input message and state element A, state element B, rotated state element C, rotated state element D, state element E, state element F, rotated state element G, and rotated state element H to generate an updated state element A, an updated state element B, an updated state element E, and an updated state element F at **510**, and commit the executed instruction to store the updated state element A, the updated state element B, the updated state element E, and the updated state element F into a location specified by the single instruction at **512**.

(78) The opcode may correspond to (e.g., identifies) the vector instruction mnemonic of VSM3RND2 (where (optional usage of) “V” represents a vector operation (e.g., VEX as discussed below), “SM3” represents the hashing standard being SM3, and “RND2” represents (e.g., only) two rounds of hashing being formed by the single instruction.

(79) In one embodiment, an instruction has the following format for its fields: VSM3RND2 SOURCE/DESTINATION, SOURCE 1, SOURCE 2// where source/destination is a first packed data source (e.g., register) and destination, source 1 is the second packed data source (e.g., register), and source 2 is the third packed data source (e.g., which may be a register or memory address). In one embodiment, each of the registers used is a (e.g., 128-bit) XMM register (e.g., XMM1, XMM2, etc.).

(80) Pseudocode for the performance of this instruction may be as follows (e.g., where dword is a doubleword size of 32-bits for each element, [x] is an element position indexed with x being from 0 to n, where n-1 is the total number of elements, << is a left rotation (that wraps around) for a given number of bits, >> is a right rotation (that wraps around) for a given number of bits, A is a logical (e.g., binary) XOR operator, | is a logical (e.g., binary) OR operator that outputs a high bit (one) if either bit or both bits in a respective bit position is set to high (one), & is a logical (e.g., binary) AND operator, and ? is a conditional operator (e.g., if condition is true ? then value X: otherwise value Y):

(81) TABLE-US-00001 // DEFINITIONS #define ROTL32(x, n) (((x) << (n)) {circumflex over ()} ((x) >> (32-(n)))) #define P0(x) (x {circumflex over ()} rot132(x, 9) {circumflex over ()} rot132(x, 17)) #define P1(x) (x {circumflex over ()} rot132(x, 15) {circumflex over ()} rot132(x, 23)) #define FF0(x, y, z) (x {circumflex over ()} y {circumflex over ()} z) #define FF1(x, y, z) (((x) & (y)) | ((x) & (z)) | ((y) & (z))) #define GG0(x, y, z) (x {circumflex over ()} y {circumflex over ()} z) #define GG1(x, y, z) (((x) & (y)) {circumflex over ()} (~(x) & (z))) #define FF(x, y, z, i) (i<16 ? FF0(x, y, z) : FF1(x, y, z)) #define GG(x, y, z, i) (i<16 ? GG0(x, y, z) : GG1(x, y, z)) // INSTRUCTION FLOW A[0] = SRC1.dword[0] B[0] = SRC1.dword[1] C[0] = SRCDEST.dword[0] D[0] = SRCDEST.dword[1] E[0] = SRC1.dword[2] F[0] = SRC1.dword[3] G[0] = SRCDEST.dword[2] H[0] = SRCDEST.dword[3] W[0] = SRC2.dword[0] W[1] = SRC2.dword[1] W[4] = SRC2.dword[2] W[5] = SRC2.dword[3] C[0] = ROTL32(C[0], 9) D[0] = ROTL32(D[0], 9) G[0] = ROTL32(G[0], 19) H[0] = ROTL32(H[0], 19) CONST = IMM<16 ? (0x79cc4519 : 0x7a879d8a); CONST = ROTL32(CONST, IMM) FOR i in 0 to 1: S1 = ROTL32((ROTL32(A[i], 12) + E[i] + CONST), 7) S2 = S1 {circumflex over ()} ROTL32(A[i], 12) T1 = FF(A[i], B[i], C[i], IMM) + D[i] + S2 + W[i] {circumflex over ()} W[i+4] T2 = GG(E[i], F[i], G[i], IMM) + H[i] + S1 + W[i] D[i+1] = C[i] C[i+1] = ROTL32(B, 9) B[i+1] = A[i] A[i+1] = T1 H[i+1] = G[i] G[i+1] = ROTL32(F[i], 19) F[i+1] = E[i] E[i+1] = P0(T2) SRCDEST.dword[0] = A[2] SRCDEST.dword[1] = B[2] SRCDEST.dword[2] = E[2] SRCDEST.dword[3] = F[2]

(82) Including such rotation operations (underlined in the above example) with a single instruction that also performs a round (e.g., two rounds) of an SM3 hashing standard thus improves the performance of a computer (e.g., processor thereof).

(83) However, in certain embodiments the rotations (underlined in the above example) for all rounds (e.g., 64 rounds) on a data block of input messages is correct except for the beginning call and the end call for the single instruction that performs this rotation (e.g., underlined in the above example) and also performs (e.g., only) two rounds of an SM3 hashing standard (e.g., SM3RND2 as discussed in reference to FIGS. 5-6). Thus, to correct this issue, a pre-rotation instruction (e.g., pre-rotate CDGH at 403 in FIG. 4B) (e.g., as discussed in reference to FIGS. 7-8) and post-rotation instruction (e.g., post-rotate CDGH at 417 in FIG. 4B) (e.g., as discussed in reference to FIGS. 9-10) are added to an instruction set architecture for a processor that supports the single instruction that performs this rotation (e.g., underlined in the above example) and also performs two rounds of

an SM3 hashing standard (e.g., SM3RND2 as discussed in reference to FIGS. 5-6).

(84) As one example, each input data block (e.g., 64 bytes of input) uses 32 calls (e.g., requests to decode/execute, etc.) of a single instruction that performs this rotation (e.g., underlined in the above example) and also performs two rounds of an SM3 hashing standard (e.g., SM3RND2 as discussed in reference to FIGS. 5-6), 13 calls of an intermediate scheduling instruction (e.g., VSM3MSG1 as discussed in reference to FIGS. 11-13 below), 13 calls of a final scheduling instruction (e.g., VSM3MSG2 as discussed in reference to FIGS. 14-16 below), one call of a pre-rotation instruction (e.g., pre-rotate CDGH at **403** in FIG. 4B) (e.g., as discussed in reference to FIGS. 7-8), and one call of a post-rotation instruction (e.g., post-rotate CDGH at **417** in FIG. 4B) (e.g., as discussed in reference to FIGS. 9-10).

(85) In one embodiment, for each input data block there are 13 calls each for VSM3MSG1 instruction and VSM3MSG2 instruction because in the first four rounds the input data message words (e.g., first 16 message words) are read directly from the input data block, for example, in an embodiment where each pair of VSM3MSG1 and VSM3MSG2 instructions calculate four words, with 68 words total, and the first 16 message words are sourced from the input data block, and the rest of the 52 words are calculated using 13 calls to a pair of VSM3MSG1 and VSM3MSG2 instructions.

(86) In one embodiment, the width of each element in the input operands is 32-bits. In the Figures herein, e.g., FIGS. 4-23, data may be loaded from a register/memory and or stored in a register or memory. In certain embodiments, the first packed data source (input) and the second packed data source (input) each have the same number of bits. In certain embodiments, the resultant packed data destination (output) has the same number of bits as the number of bits in the first packed data source (input) or the second packed data source (input).

(87) In certain embodiments, a VSM3RND2 instruction (e.g., its opcode) indicates the rotational values above (e.g., 12 for functions S1 and S2; 9 for underlined functions C[i] and D[i], 19 for underlined functions G[i] and H[i], 9 for function C[i+1]; 19 for function G[i+1]; 9 and 17 for function P0; and 15 and 23 for function P1). This improves the functioning of a computer as these values are thus not to be supplied as operands (e.g., and thus it will not utilize the time and energy taken to load those operands) to those embodiments of the instruction.

(88) In certain embodiments, some or all of the data may be accessed in (e.g., system) memory. The input and output vector values and sizes herein are also examples, and other values and sizes may be utilized. The data may be according to big-endian or little-endian order.

(89) FIG. 6 illustrates circuitry **600** including an execution circuit **610** with rotation circuitry **612** and SM3 hashing round circuitry **614** according to embodiments of the disclosure. In certain embodiments, decode circuit (e.g., decode circuit **102** in FIG. 1 or decode circuit **206** in FIG. 2) decodes an instruction into a decoded instruction that causes execution circuit **610** to perform a rotation operation with rotation circuitry **612** and SM3 hashing round circuitry **614** (for example, the decoded instruction indicates to execution circuit **610** which components to use, e.g., here to use the rotation circuitry **612** and SM3 hashing round circuitry **614**, the direction(s) of rotation, and the values of the rotation constants and hashing constants). In the depicted embodiment, an instruction format may include one or more fields that identifies the first packed data destination/source **601**, second packed data source **602**, and the third packed data source **603**. In the depicted embodiment, the first packed data destination/source **601**, second packed data source **602**, and the third packed data source **603** each include four elements (indexed 0-3). However, it should be understood that any number of elements, or a proper subset of any numbers, may be utilized with the instruction(s) disclosed herein. In certain embodiments, each element in a packed data source has a same bit width (e.g., a byte or bytes) and/or each of the elements in each of the other packed data sources.

(90) In one embodiment, first packed data destination/source **601** is a four element location that stores state elements C, D, G, and H, respectively, second packed data source **602** is a four element

location that stores state elements A, B, E, and F, respectively, and the third packed data source **603** is a four element location that stores input message elements WM, W[i+1], W[i+4], and W[i+5], respectively for each round “i”. Input message elements may be sourced from a scheduling instruction(s).

(91) In certain embodiments, rotation circuitry **612** of execution circuit **610** performs a rotation operation on each element such that element C from the first element of first packed data source **601** is left rotated nine bits at 32-bit granularity, element D from the second element of first packed data source **601** is left rotated nine bits at 32-bit granularity, element G from the third element of first packed data source **601** is left rotated nineteen bits at 32-bit granularity, and element H from the fourth element of first packed data source **601** is left rotated nineteen bits at 32-bit granularity as indicated by the instruction (e.g., by the opcode). In one embodiment, the rotation operations are performed on each element in parallel.

(92) The execution circuit **610** then causes the results of the rotations from rotation circuitry **612** to be sent to SM3 hashing round circuitry **614**, for example, along with the other operands (e.g., with the other operands bypassing the rotation circuitry **612**). In certain embodiments, the SM3 hashing round circuitry **614** performs hashing operations (e.g., two rounds as discussed above) on the results of the rotations from rotation circuitry **612** and the state elements A, B, E, and F from the second packed data source **602** and the input message elements WM, W[i+1], W[i+4], and W[i+5] from the third packed data source **603** to generate a hash two round result (e.g., including updated state elements A, B, E, and F). The packed data hashing result (e.g., updated state elements A, B, E, and F) may then be stored (e.g., committed to be visible by other instructions) into destination **601**.

(93) FIG. 7 illustrates a method **700** of processing a pre-rotation instruction according to a SM3 hashing standard according to embodiments of the disclosure. A processor (e.g., or processor core) may perform method **700**, e.g., in response to receiving a request to execute an instruction from software. Depicted method **700** includes processing a SM3 hashing standard instruction by: fetch a single instruction having an opcode that indicates a pre-rotation according to a SM3 hashing standard is to be performed and at least one field that identifies four (e.g., 32-bit) initial state elements C, D, G, and H for an initial round of a data block (e.g., for input into single instruction VSM3RNDS2 in the format that is assumed to be correct) according to the SM3 hashing standard (e.g., and a location to write a resultant) at **702**, decode the instruction into a decoded instruction at **704**, retrieve data associated with the identified source operands at **706**, (optionally) schedule the decoded instruction for execution at **708**, execute the decoded instruction to rotate initial state element C right by 9 bits to form a pre-rotated state element C, rotate initial state element D right by 9 bits to form a pre-rotated state element D, rotate initial state element G right by 19 bits to form a pre-rotated state element G, rotate initial state element H right by 19 bits to form a pre-rotated state element H at **710**, and commit the executed instruction to store the pre-rotated state element C, pre-rotated state element D, pre-rotated state element G, and pre-rotated state element H into a location specified by the single instruction at **712**.

(94) The opcode may correspond to (e.g., identifies) the vector instruction mnemonic of PREROTATESM3 (e.g., where “PREROTATESM3” represents pre-rotating (e.g., reverse rotating) state elements C, D, G, and G according to a SM3 hashing standard).

(95) In one embodiment, an instruction has the following format for its fields: PREROTATESM3 SOURCE/DESTINATION// where source/destination is a first packed data source (e.g., register) and destination. In one embodiment, each the source/destination is a register (e.g., a XMM register).

(96) Thus, such an embodiment of a pre-rotation instruction provides data that, when input into single instruction that performs this rotation (e.g., underlined in the above example) and also performs two rounds of an SM3 hashing standard (e.g., SM3RNDS2 as discussed in reference to FIGS. 5-6), produces a correct output for a first round of a plurality of rounds for an input message. In one embodiment, the width of each element in the input operand is 32-bits. In the Figures herein,

e.g., FIGS. 4-23, data may be loaded from a register/memory and or stored in a register or memory. In certain embodiments, the first packed data source (input) and the second packed data source (input) each have the same number of bits. In certain embodiments, the resultant packed data destination (output) has the same number of bits as the number of bits in the first packed data source (input) or the second packed data source (input).

(97) In certain embodiments, some or all of the data may be accessed in (e.g., system) memory. The input and output vector values and sizes herein are also examples, and other values and sizes may be utilized. The data may be according to big-endian or little-endian order.

(98) FIG. 8 illustrates circuitry 800 including an execution circuit 810 with rotation circuitry 812 according to embodiments of the disclosure. In certain embodiments, decode circuit (e.g., decode circuit 102 in FIG. 1 or decode circuit 206 in FIG. 2) decodes an instruction into a decoded instruction that causes execution circuit 810 to perform a pre-rotation operation with rotation circuitry 812 for an initial SM3 hashing round (for example, the decoded instruction indicates to execution circuit 810 which components to use, e.g., here to use the rotation circuitry 812, the direction of rotation(s), and the values of the rotation constants). In the depicted embodiment, an instruction format may include one or more fields that identifies the first packed data destination/source 801. In the depicted embodiment, the first packed data destination/source 801 includes four elements (indexed 0-3). However, it should be understood that any number of elements, or a proper subset of any numbers, may be utilized with the instruction(s) disclosed herein. In certain embodiments, each element in a packed data source has a same bit width (e.g., a byte or bytes) and/or each of the elements in each of the other packed data sources.

(99) In one embodiment, first packed data destination/source 801 is a four element location that stores state elements C, D, G, and H, respectively. Input message elements may be sourced from a scheduling instruction(s).

(100) In certain embodiments, rotation circuitry 812 of execution circuit 810 performs a pre-rotation operation on each element such that element C from the first element of first packed data source 801 is right rotated nine bits at 32-bit granularity, element D from the second element of first packed data source 801 is right rotated nine bits at 32-bit granularity, element G from the third element of first packed data source 801 is right rotated nineteen bits at 32-bit granularity, and element H from the fourth element of first packed data source 801 is right rotated nineteen bits at 32-bit granularity as indicated by the instruction (e.g., by the opcode). In one embodiment, the pre-rotation operations are performed on each element in parallel.

(101) The packed data pre-rotation result (e.g., pre-rotated state elements C, D, G, and H) may then be stored (for example, committed to be visible by other instructions, e.g., which include a SM3RNDs2 instruction) into destination 801.

(102) FIG. 9 illustrates a method 900 of processing a post-rotation instruction according to a SM3 hashing standard according to embodiments of the disclosure. A processor (e.g., or processor core) may perform method 900, e.g., in response to receiving a request to execute an instruction from software. Depicted method 900 includes processing a SM3 hashing standard instruction by: fetch a single instruction having an opcode that indicates a post-rotation according to a SM3 hashing standard is to be performed and at least one field that identifies four (e.g., 32-bit) updated state elements C, D, G, and H for a final round of a data block (e.g., C, D, G, and H output for a final round from a single instruction VSM3RNDs2) according to the SM3 hashing standard (e.g., and a location to write a resultant) at 902, decode the instruction into a decoded instruction at 904, retrieve data associated with the identified source operands at 906, (optionally) schedule the decoded instruction for execution at 908, execute the decoded instruction to rotate updated state element C left by 9 bits to form a final state element C of the data block, rotate updated state element D left by 9 bits to form a final state element D of the data block, rotate updated state element G left by 19 bits to form a final state element G of the data block, rotate updated state element H left by 19 bits to form a final state element H of the data block at 910, and commit the

executed instruction to store the final state element C, final state element D, final state element G, and final state element H into a location specified by the single instruction at **912**.

(103) The opcode may correspond to (e.g., identifies) the vector instruction mnemonic of POSTROTATESM3 (e.g., where “POSTROTATESM3” represents post-rotating state elements C, D, G, and G according to a SM3 hashing standard).

(104) In one embodiment, an instruction has the following format for its fields:

POSTROTATESM3 SOURCE/DESTINATION// where source/destination is a first packed data source (e.g., register) and destination. In one embodiment, each the source/destination is a register (e.g., a XMM register).

(105) Thus, such an embodiment of a post-rotation instruction provides data that, when taking output from a single instruction that performs this rotation (e.g., underlined in the above example) and also performs two rounds of an SM3 hashing standard (e.g., SM3RND2 as discussed in reference to FIGS. 5-6), produces a correct output for a last round of a plurality of rounds for an input message. In one embodiment, the width of each element in the input operand is 32-bits. In the Figures herein, e.g., FIGS. 4-23, data may be loaded from a register/memory and or stored in a register or memory. In certain embodiments, the first packed data source (input) and the second packed data source (input) each have the same number of bits. In certain embodiments, the resultant packed data destination (output) has the same number of bits as the number of bits in the first packed data source (input) or the second packed data source (input).

(106) In certain embodiments, some or all of the data may be accessed in (e.g., system) memory. The input and output vector values and sizes herein are also examples, and other values and sizes may be utilized. The data may be according to big-endian or little-endian order.

(107) FIG. 10 illustrates circuitry **1000** including an execution circuit **1010** with rotation circuitry **1012** according to embodiments of the disclosure. In certain embodiments, decode circuit (e.g., decode circuit **102** in FIG. 1 or decode circuit **206** in FIG. 2) decodes an instruction into a decoded instruction that causes execution circuit **1010** to perform a post-rotation operation with rotation circuitry **1012** for a final SM3 hashing round (for example, the decoded instruction indicates to execution circuit **1010** which components to use, e.g., here to use the rotation circuitry **1012**, the direction of rotation(s), and the values of the rotation constants). In the depicted embodiment, an instruction format may include one or more fields that identifies the first packed data destination/source **1001**. In the depicted embodiment, the first packed data destination/source **1001** includes four elements (indexed 0-3). However, it should be understood that any number of elements, or a proper subset of any numbers, may be utilized with the instruction(s) disclosed herein. In certain embodiments, each element in a packed data source has a same bit width (e.g., a byte or bytes) and/or each of the elements in each of the other packed data sources.

(108) In one embodiment, first packed data destination/source **1001** is a four element location that stores state elements C, D, G, and H, respectively. Input message elements may be sourced from a SM3RND2 instruction.

(109) In certain embodiments, rotation circuitry **1012** of execution circuit **1010** performs a post-rotation operation on each element such that element C from the first element of first packed data source **1001** is left rotated nine bits at 32-bit granularity, element D from the second element of first packed data source **1001** is left rotated nine bits at 32-bit granularity, element G from the third element of first packed data source **1001** is left rotated nineteen bits at 32-bit granularity, and element H from the fourth element of first packed data source **1001** is left rotated nineteen bits at 32-bit granularity as indicated by the instruction (e.g., by the opcode). In one embodiment, the post-rotation operations are performed on each element in parallel.

(110) The packed data post-rotation result (e.g., post-rotated state elements C, D, G, and H) may then be stored (for example, committed to be visible by other instructions) into destination **1001**.

(111) As discussed in reference to FIGS. 4A and 4B, certain embodiments herein utilize scheduling instruction(s) to generate input message elements (e.g., for consumption by a SM3RND2

instruction). FIGS. **11-13** depict embodiments for an intermediate message instruction, and FIGS. **14-16** depict embodiments for a final message instruction.

(112) FIG. **11** illustrates a method **1100** of processing an intermediate message instruction according to a SM3 hashing standard according to embodiments of the disclosure. A processor (e.g., or processor core) may perform method **1100**, e.g., in response to receiving a request to execute an instruction from software. Depicted method **1100** includes processing a SM3 hashing standard instruction by: fetch a single instruction having an opcode that indicates an intermediate calculation for an input message according to a SM3 hashing standard is to be performed and at least one first field that identifies eleven (e.g., 32-bit) message elements $W[i]$, $W[i+1]$, $W[i+2]$, $W[i+3]$, $W[i+7]$, $W[i+8]$, $W[i+9]$, $W[i+10]$, $W[i+13]$, $W[i+14]$, and $W[i+15]$ for a round “i” (e.g., an immediately previous round) according to the SM3 hashing standard (e.g., and a location to write a resultant) at **1102**, decode the instruction into a decoded instruction at **1104**, retrieve data associated with the identified source operands at **1106**, (optionally) schedule the decoded instruction for execution at **1108**, execute the decoded instruction to determine four (e.g., 32-bit) intermediate message elements (e.g., $X[3]$, $X[2]$, $X[1]$, and $X[0]$ in FIG. **13**) from the eleven (e.g., 32-bit) message elements $W[i]$, $W[i+1]$, $W[i+2]$, $W[i+3]$, $W[i+7]$, $W[i+8]$, $W[i+9]$, $W[i+10]$, $W[i+13]$, $W[i+14]$, and $W[i+15]$ at **1110**, and commit the executed instruction to store the four intermediate message elements into a location specified by the single instruction at **1112**.

(113) The opcode may correspond to (e.g., identifies) the vector instruction mnemonic of VSM3MSG1 (e.g., where (optional usage of) “V” represents a vector operation (e.g., VEX as discussed below), “SM3” represents the hashing standard being SM3, and “MSG1” represents an intermediate message “1” used to generate a final message “2”, with the data from that final message used as an input of a SM3RND2 instruction).

(114) In one embodiment, an instruction has the following format for its fields: VSM3MSG1 SOURCE/DESTINATION, SOURCE 1, SOURCE 2// where source/destination is a first packed data source (e.g., register) and destination, source 1 is the second packed data source (e.g., register), and source 2 is the third packed data source (e.g., which may be a register or memory address). In one embodiment, each of the registers used is a (e.g., 128-bit) XMM register (e.g., XMM1, XMM2, etc.).

(115) Pseudocode for the performance of this instruction may be as follows (e.g., where dword is a doubleword size of 32-bits for each element, $[x]$ is an element position indexed with x being from 0 to n , where $n-1$ is the total number of elements, $\ll$ is a left rotation (that wraps around) for a given number of bits, $\gg$ is a right rotation (that wraps around) for a given number of bits, $\oplus$ is a logical (e.g., binary) XOR operator, $\mid$ is a logical (e.g., binary) OR operator that outputs a high bit (one) if either bit or both bits in a respective bit position is set to high (one), $\&$ is a logical (e.g., binary) AND operator, and $?$ is a conditional operator (e.g., if condition is true ? then value X: otherwise value Y):

(116) TABLE-US-00002 // DEFINITIONS #define ROTL32(x, n) (((x) << (n)) {circumflex over ()} ((x) >> (32-(n)))) #define P1(x) (x {circumflex over ()} rot132(x, 15) {circumflex over ()} rot132(x, 23)) // INSTRUCTION FLOW: $W[7] = \text{SRCDST.dword}[0]$ $W[8] = \text{SRCDST.dword}[1]$ $W[9] = \text{SRCDST.dword}[2]$ $W[10] = \text{SRCDST.dword}[3]$ $W[13] = \text{SRC1.dword}[0]$ $W[14] = \text{SRC1.dword}[1]$ $W[15] = \text{SRC1.dword}[2]$ $W[0] = \text{SRC2.dword}[0]$ $W[1] = \text{SRC2.dword}[1]$ $W[2] = \text{SRC2.dword}[2]$ $W[3] = \text{SRC2.dword}[3]$ $\text{TMP3} = W[10] \{ \text{circumflex over ()} W[3] \}$ $\text{TMP2} = W[9] \{ \text{circumflex over ()} W[2] \}$ $\text{ROTL32}(W[15], 15)$ $\text{TMP1} = W[8] \{ \text{circumflex over ()} W[1] \}$ $\text{ROTL32}(W[14], 15)$ $\text{TMP0} = W[7] \{ \text{circumflex over ()} W[0] \}$ $\text{ROTL32}(W[13], 15)$ $\text{SRCDST.dword}[0] = \text{P1}(\text{TMP0})$ $\text{SRCDST.dword}[1] = \text{P1}(\text{TMP1})$ $\text{SRCDST.dword}[2] = \text{P1}(\text{TMP2})$ $\text{SRCDST.dword}[3] = \text{P1}(\text{TMP3})$

(117) FIG. **12** illustrates circuitry **1200** including an execution circuit **1210** with intermediate SM3 message calculation circuitry **1212** according to embodiments of the disclosure. In certain

embodiments, decode circuit (e.g., decode circuit **102** in FIG. **1** or decode circuit **206** in FIG. **2**) decodes an instruction into a decoded instruction that causes execution circuit **1210** to perform a SM3 message calculation operation with intermediate SM3 message calculation circuitry **1212** (for example, the decoded instruction indicate to execution circuit **1210** which components to use, e.g., here to use the intermediate SM3 message calculation circuitry **1212**). In the depicted embodiment, an instruction format may include one or more fields that identifies the first packed data destination/source **1201**, second packed data source **1202**, and the third packed data source **1203**. In the depicted embodiment, the first packed data destination/source **1201**, second packed data source **1202**, and the third packed data source **1203** each include four elements (indexed 0-3). However, it should be understood that any number of elements, or a proper subset of any numbers, may be utilized with the instruction(s) disclosed herein. In certain embodiments, each element in a packed data source has a same bit width (e.g., a byte or bytes) and/or each of the elements in each of the other packed data sources.

(118) In one embodiment, first packed data destination/source **1201** is a four element location that stores input message elements $W[i+7]$, $W[i+8]$, $W[i+9]$, and $W[i+10]$, respectively, second packed data source **1202** is a three (e.g., or four, with one element not utilized) element location that stores input message elements $W[i+13]$, $W[i+14]$, and $W[i+15]$, respectively, and the third packed data source **1203** is a four element location that stores input message elements $W[i+0]$, $W[i+1]$, $W[i+2]$, and $W[i+3]$, respectively. Input message elements may be sourced from an input message, e.g., a hashing input. The input message may be from a previous message (e.g., a previous round to the round the current message is being determined).

(119) In certain embodiments, intermediate SM3 message calculation circuitry **1212** of execution circuit **1210** performs message generation operations on the input message elements to generate a plurality (e.g., four) intermediate message elements as indicated by the instruction (e.g., by the opcode). In one embodiment, the message generation operations are performed on each element in parallel. In one embodiment, the intermediate SM3 message calculation circuitry **1212** is an instance of the intermediate SM3 message calculation circuitry **1300** in FIG. **13**. The packed data intermediate message result (e.g., $X[3]$, $X[2]$, $X[1]$, and $X[0]$) may then be stored (e.g., committed to be visible by other instructions) into destination **1201**, for example, for use by a next (e.g., in program order) of the pair of message scheduling instructions.

(120) FIG. **13** illustrates intermediate SM3 message calculation circuitry **1300** according to embodiments of the disclosure. Note that the element reference format of $W[x]$ may be the same as $W.sub.i+x$ in FIG. **13**, where X is an integer and i is the round number. In one embodiment, first packed data destination/source **1301** is a four element location (e.g., register) that stores input message elements $W[i+7]$, $W[i+8]$, $W[i+9]$, and $W[i+10]$, respectively, second packed data source **1202** is a three (e.g., or four, with one element not utilized) element location (e.g., register) that stores input message elements $W[i+13]$, $W[i+14]$, and $W[i+15]$, respectively, and the third packed data source **1303** is a four element location (e.g., register) that stores input message elements $W[i+0]$, $W[i+1]$, $W[i+2]$, and $W[i+3]$, respectively. Input message elements may be sourced from an input message, e.g., a hashing input. The input message may be from a previous message (e.g., a previous round to the round the current message is being determined for).

(121) In the depicted embodiment, the input message elements are routed according to the couplings (e.g., wires) as shown to: the left (e.g., 15 bits to the left with a 32 bit element size) shifter circuits **1304B-1304D**, XOR logic gate circuits **1306A-1306D**, and permutation function circuits **1308A-1308D** (e.g., to perform the permutation function in function 5 above). The output of four intermediate message elements (e.g., $X[0]$, $X[1]$, $X[2]$, and $X[3]$) from the intermediate SM3 message calculation circuitry **1300** may then be saved to first packed data destination/source **1301**. A logic gate circuit may be implemented with positive or negative logic gates.

(122) FIG. **14** illustrates a method **1400** of processing a final message instruction according to a SM3 hashing standard according to embodiments of the disclosure. Final message instruction may

be next (e.g., in program order) to an intermediate message instruction that is providing its output as an input to the final message instruction. A processor (e.g., or processor core) may perform method **1400**, e.g., in response to receiving a request to execute an instruction from software. Depicted method **1400** includes processing a SM3 hashing standard instruction by: fetch a single instruction having an opcode that indicates a final calculation for an input message according to a SM3 hashing standard is to be performed and at least one first field that identifies twelve (e.g., 32-bit) message elements $W[i+3]$, $W[i+4]$, $W[i+5]$, $W[i+6]$, $W[i+10]$, $W[i+11]$, $W[i+12]$, $W[i+13]$, and four intermediate message elements from an intermediate calculation (e.g., from an instance of instruction VSM3MSG1) for a round according to the SM3 hashing standard (e.g., and a location to write a resultant) at **1402**, decode the instruction into a decoded instruction at **1404**, retrieve data associated with the identified source operands at **1406**, (optionally) schedule the decoded instruction for execution at **1408**, execute the decoded instruction to determine four (e.g., 32-bit) final message elements from the twelve (e.g., 32-bit) message elements $W[i+3]$, $W[i+4]$, $W[i+5]$, $W[i+6]$, $W[i+10]$, $W[i+11]$, $W[i+12]$, $W[i+13]$, and the four intermediate message elements from the intermediate calculation at **1410**, and commit the executed instruction to store the four final message elements into a location specified by the single instruction at **1412**.

(123) The opcode may correspond to (e.g., identifies) the vector instruction mnemonic of VSM3MSG2 (e.g., where (optional usage of) “V” represents a vector operation (e.g., VEX as discussed below), “SM3” represents the hashing standard being SM3, and “MSG1” represents a final message “2” generated from an intermediate message “1”, with the data from the final message used as an input of a SM3RND2 instruction).

(124) In one embodiment, an instruction has the following format for its fields: VSM3MSG2 SOURCE/DESTINATION, SOURCE 1, SOURCE 2// where source/destination is a first packed data source (e.g., register) and destination, source 1 is the second packed data source (e.g., register), and source 2 is the third packed data source (e.g., which may be a register or memory address). In one embodiment, each of the registers used is a (e.g., 128-bit) XMM register (e.g., XMM1, XMM2, etc.).

(125) Pseudocode for the performance of this instruction may be as follows (e.g., where dword is a doubleword size of 32-bits for each element, $[x]$ is an element position indexed with x being from 0 to n , where $n-1$ is the total number of elements, $\ll$ is a left rotation (that wraps around) for a given number of bits, $\gg$ is a right rotation (that wraps around) for a given number of bits, $\wedge$ is a logical (e.g., binary) XOR operator, $\vee$ is a logical (e.g., binary) OR operator that outputs a high bit (one) if either bit or both bits in a respective bit position is set to high (one), $\&$ is a logical (e.g., binary) AND operator, and $?$ is a conditional operator (e.g., if condition is true ? then value X: otherwise value Y):

(126) TABLE-US-00003 // DEFINITIONS #define ROTL32(x, n) (((x) << (n)) {circumflex over ()} ((x) >> (32-(n)))) // INSTRUCTION FLOW WTMP[0] = SRCDST.dword[0] WTMP[1] = SRCDST.dword[1] WTMP[2] = SRCDST.dword[2] WTMP[3] = SRCDST.dword[3] $W[3] = \text{SRC1.dword}[0]$ $W[4] = \text{SRC1.dword}[1]$ $W[5] = \text{SRC1.dword}[2]$ $W[6] = \text{SRC1.dword}[3]$ $W[10] = \text{SRC2.dword}[0]$ $W[11] = \text{SRC2.dword}[1]$ $W[12] = \text{SRC2.dword}[2]$ $W[13] = \text{SRC2.dword}[3]$ $W[16] = (\text{ROTL32}(W[3], 7) \{circumflex over ()\} W[10] \{circumflex over ()\} WTMP[0] W[17] = (\text{ROTL32}(W[4], 7) \{circumflex over ()\} W[11] \{circumflex over ()\} WTMP[1] W[18] = (\text{ROTL32}(W[5], 7) \{circumflex over ()\} W[12] \{circumflex over ()\} WTMP[2] W[19] = (\text{ROTL32}(W[6], 7) \{circumflex over ()\} W[13] \{circumflex over ()\} WTMP[3] W[19] = (W[19] \{circumflex over ()\} \text{ROTL32}(W[16], 6) \{circumflex over ()\} \text{ROTL32}(W[16], 15) \{circumflex over ()\} \text{ROTL32}(W[16], 30))$ SRCDST.dword[0] = $W[16]$ SRCDST.dword[1] = $W[17]$ SRCDST.dword[2] = $W[18]$ SRCDST.dword[3] = $W[19]$

(127) FIG. 15 illustrates circuitry **1500** including an execution circuit **1510** with final SM3 message calculation circuitry **1512** according to embodiments of the disclosure. In certain embodiments, decode circuit (e.g., decode circuit **102** in FIG. 1 or decode circuit **206** in FIG. 2) decodes an

instruction into a decoded instruction that causes execution circuit **1510** to perform a SM3 message calculation operation with final SM3 message calculation circuitry **1512** (for example, the decoded instruction indicate to execution circuit **1510** which components to use, e.g., here to use the final SM3 message calculation circuitry **1512**). In the depicted embodiment, an instruction format may include one or more fields that identifies the first packed data destination/source **1501**, second packed data source **1502**, and the third packed data source **1503**. In the depicted embodiment, the first packed data destination/source **1501**, second packed data source **1502**, and the third packed data source **1503** each include four elements (indexed 0-3). However, it should be understood that any number of elements, or a proper subset of any numbers, may be utilized with the instruction(s) disclosed herein. In certain embodiments, each element in a packed data source has a same bit width (e.g., a byte or bytes) and/or each of the elements in each of the other packed data sources.

(128) In one embodiment, first packed data destination/source **1501** is a four element location that stores intermediate message elements (e.g., $X[0]$, $X[1]$, $X[2]$, and $X[3]$) from an intermediate SM3 message calculation instruction (e.g., VSM3MSG1), respectively, second packed data source **1502** is a four element location that stores input message elements $W[i+3]$, $W[i+4]$, $W[i+5]$, and $W[i+6]$, respectively, and the third packed data source **1503** is a four element location that stores input message elements $W[i+10]$, $W[i+11]$, $W[i+12]$, and $W[i+13]$, respectively. Input message elements may be sourced from an input message, e.g., a hashing input, and a previous intermediate SM3 message calculation instruction (e.g., VSM3MSG1). The input message may be from a previous message (e.g., a previous round to the round the current message is being determined).

(129) In certain embodiments, final SM3 message calculation circuitry **1512** of execution circuit **1510** performs message generation operations on the input message elements and intermediate message elements to generate a plurality (e.g., four) final message elements as indicated by the instruction (e.g., by the opcode). In one embodiment, the message generation operations are performed on each element in parallel. In one embodiment, the final SM3 message calculation circuitry **1512** is an instance of the final SM3 message calculation circuitry **1600** in FIG. 16. The packed data final message result (e.g., $W[i+16]$, $W[i+17]$, $W[i+18]$, and $W[i+19]$) may then be stored (e.g., committed to be visible by other instructions) into destination **1501**, for example, for use by a next (e.g., in program order) of the pair of message scheduling instructions.

(130) FIG. 16 illustrates final SM3 message calculation circuitry **1600** according to embodiments of the disclosure. Note that the element reference format of $W[x]$ may be the same as $W.sub.i+x$ in FIG. 16, where X is an integer and i is the round number. In one embodiment, first packed data destination/source **1601** is a four element location (e.g., register) that stores intermediate message elements (e.g., $X[0]$, $X[1]$, $X[2]$, and $X[3]$ from an intermediate SM3 message calculation instruction), respectively, second packed data source **1202** is a four element location (e.g., register) that stores input message elements $W[i+3]$, $W[i+4]$, $W[i+5]$, and $W[i+6]$, respectively, and the third packed data source **1603** is a four element location (e.g., register) that stores input message elements $W[i+10]$, $W[i+11]$, $W[i+12]$, and $W[i+13]$, respectively. Input message elements may be sourced from an input message, e.g., a hashing input. Input message elements may be sourced from an input message, e.g., a hashing input, and a previous intermediate SM3 message calculation instruction (e.g., VSM3MSG1). The input message may be from a previous message (e.g., a previous round to the round the current message is being determined).

(131) In the depicted embodiment, the input message elements are routed according to the couplings (e.g., wires) as shown to: the left (e.g., 7 bits to the left with a 32 bit element size) shifter circuits **1604B-1304D**, left (e.g., 6 bits to the left with a 32 bit element size) shifter circuit **1608A**, left (e.g., 15 bits to the left with a 32 bit element size) shifter circuit **1608B**, left (e.g., 30 bits to the left with a 32 bit element size) shifter circuit **1608C**, XOR logic gate circuit **1610**, and XOR logic gate circuit **1612**. The output of four final message elements $W[i+16]$, $W[i+17]$, $W[i+18]$, and $W[i+19]$ from the final SM3 message calculation circuitry **1600** may then be saved to first packed data destination/source **1601** (e.g., for use as an input of a SM3RND2 instruction). A logic gate

circuit may be implemented with positive or negative logic gates.

(132) In certain embodiments, an execution circuit includes circuitry for both SM3 hashing standards and SHA512 hashing standards.

(133) Instructions for a SHA512 Hashing Standard

(134) FIG. 17 illustrates SHA512 hashing circuitry 1700 according to embodiments of the disclosure. Circuitry for each of $Ch(x, y, z)$, $Maj(x, y, z)$, $\Sigma.sub.0$, $\Sigma.sub.1$ (e.g., as defined below) implement (e.g., as execution circuitry for a SHA512 hashing instruction) respective compression functions that are defined according to the SHA512 (i.e., SHA-512) specification, and each state (e.g., state vector A-H) is initiated with a respective 64-bit (e.g., and not a 32-bit) initial value, and W_i , $i=0$ to 79 (e.g., and not 63), are (e.g., 64-bit) values derived from a (e.g., 1024-bit) message. K_t in constant storage 1706 represents a constant value to be used for the iteration t of the hash computation. In the below, “ $\{ \text{circumflex over ()} \}$ ” is the bitwise AND operation, “ $\oplus$ ” is the bitwise XOR (“exclusive-OR”) operation, “ $\neg$ ” is the bitwise complement operation (e.g., it changes 1 to 0 and 0 to 1), $ROTR.sup.n$ indicates to rotate right by “ n ” number of bits (e.g., circular rotate with wrapping around), and $SHR.sup.n$ indicates to shift right by “ n ” number of bits (e.g., shift without wrapping around). The circuits depicted as a square with a plus in it are adder circuits that each add their respective addends modulo of some particular value (e.g., modulo $2.sup.32$ or $2.sup.64$) to produce their respective sums.

$Ch(x,y,z)=(x\{ \text{circumflex over ()} \}y)\oplus(\neg x\{ \text{circumflex over ()} \}z)$ Function 7

$Maj(x,y,z)=(x\{ \text{circumflex over ()} \}y)\oplus(x\{ \text{circumflex over ()} \}z)\oplus(y\{ \text{circumflex over ()} \}z)$
Function 8

$\Sigma.sub.0(x)=ROTR.sup.28(x)\oplus ROTR.sup.34(x)\oplus ROTR.sup.39(x)$ Function 9

$\Sigma.sub.1(x)=ROTR.sup.14(x)\oplus ROTR.sup.18(x)\oplus ROTR.sup.41(x)$ Function 10

$\sigma.sub.0=ROTR.sup.1(x)\oplus ROTR.sup.8(x)\oplus SHR.sup.7(x)$ Function 11

$\sigma.sub.1(x)=ROTR.sup.19(x)\oplus ROTR.sup.61(x)\oplus SHR.sup.6(x)$ Function 12

(135) In certain embodiments, intermediate message instruction (e.g., VSHA512MSG1) according to a SHA512 hashing standard and final message instruction (e.g., VSHA512MSG2) according to a SHA512 hashing standard are decoded and executed (e.g., at 1702) to provide input messages (e.g., for one or two rounds of hashing), for example, and storing the input messages into storage 1704. In certain embodiments, execution circuitry 1708 generates (e.g., via decoding and executing of VSHA512RNDS2 instruction at 1710) two (or one in another embodiment) rounds of updated state elements (e.g., updated state elements A, B, E, and F). Updated state elements may then be used by hashing circuitry 1700 to produce (e.g., 512-bits) the SHA512 message digest (e.g., signature) (e.g., the message digest generated after 80 rounds of hashing on an input block).

(136) In certain embodiments, a VSHA512RNDS2 instruction (e.g., its opcode) indicates the rotational and shift values above (e.g., 28, 34, and 39 for function 9; 14, 18, and 41 for function 10; 1, 8, and 7 for function 11; and 19, 61, and 6 for function 12). This improves the functioning of a computer as these values are thus not to be supplied as operands (e.g., and thus it will not utilize the time and energy taken to load those operands) to those embodiments of the instruction.

(137) In one embodiment, a SHA384 hashing standard is the same as a SHA-512 hashing standard except that: the initial hash values h_0 through h_7 are different (e.g., taken from the 9th through 16th primes), and the output is constructed by omitting h_6 and h_7 .

(138) FIG. 18 illustrates a method 1800 of processing a hashing (e.g., round) instruction according to a SHA512 hashing standard according to embodiments of the disclosure. A processor (e.g., or processor core) may perform method 1800, e.g., in response to receiving a request to execute an instruction from software. Depicted method 1800 includes processing a SHA512 hashing standard instruction by: fetch a single instruction having an opcode that indicates two rounds of hashing according to a SHA512 hashing standard are to be performed and at least one first field that identifies eight (e.g., 64-bit) state elements (e.g., working variables) A, B, C, D, E, F, G, and H for

a round according to the SHA512 hashing standard and at least one second field that identifies an input message (e.g., of the message schedule) (e.g., and a field that indicates a location to write a resultant) at **1802**, decode the instruction into a decoded instruction at **1804**, retrieve data associated with the identified source operands at **1806**, (optionally) schedule the decoded instruction for execution at **1808**, execute the decoded instruction to perform two rounds according to the SHA512 hashing standard on the input message and state elements A, B, C, D, E, F, G, and H to generate an updated state element A, an updated state element B, an updated state element E, and an updated state element F at **1810**, and commit the executed instruction to store the updated state element A, the updated state element B, the updated state element E, and the updated state element F into a location specified by the single instruction at **1812**.

(139) The opcode may correspond to (e.g., identifies) the vector instruction mnemonic of VSHA512RNDS2 (e.g., where (optional usage of) “V” represents a vector operation (e.g., VEX as discussed below), “SHA512” represents the hashing standard being SHA512, and “RNDS2” represents (e.g., only) two rounds of hashing being formed by the single instruction).

(140) In one embodiment, an instruction has the following format for its fields: VSHA512RNDS2 SOURCE/DESTINATION, SOURCE 1, SOURCE 2// where source/destination is a first packed data source (e.g., register) and destination, source 1 is the second packed data source (e.g., register), and source 2 is the third packed data source (e.g., which may be a register or memory address). In one embodiment, each of the registers used is a (e.g., 256-bit) YMM register (e.g., YMM1, YMM2, etc.). A memory address may be a reference to 256-bit of data, e.g., referred to as “m256”.

(141) Pseudocode for the performance of this instruction may be as follows (e.g., where qword is a quadword size of 64-bits for each element, [x] is an element position indexed with x being from 0 to n, where n-1 is the total number of elements, << is a left rotation (that wraps around) for a given number of bits, >> is a right rotation (that wraps around) for a given number of bits, {circumflex over ()} is a logical (e.g., binary) XOR operator, | is a logical (e.g., binary) OR operator that outputs a high bit (one) if either bit or both bits in a respective bit position is set to high (one), & is a logical (e.g., binary) AND operator, and ? is a conditional operator (e.g., if condition is true ? then value X: otherwise value Y):

(142) TABLE-US-00004 // DEFINITIONS #define ROL64(x, n) (((x) << (n)) {circumflex over ()} ((x) >> (64-(n)))) #define ROR64(x, n) (((x) >> (n)) {circumflex over ()} ((x) << (64-(n)))) #define SHR64(x, n) ((x) >> (n)) #define S0(a) (ROR64(a,28) {circumflex over ()} ROR64(a, 34) {circumflex over ()} ROR64(a,39)) #define S1(e) (ROR64(e,14) {circumflex over ()} ROR64(e, 18) {circumflex over ()} ROR64(e,41)) #define MAJ(a,b,c) (((a) & (b)) {circumflex over ()} ((a) & (c)) {circumflex over ()} ((b) & (c))) #define CH(e,f,g) (((e) & (f)) {circumflex over ()} ((g) & ~(e))) // INSTRUCTION FLOW A[0] := SRC1.qword[3] B[0] := SRC1.qword[2] C[0] := SRCDST.qword[3] D[0] := SRCDST.qword[2] E[0] := SRC1.qword[1] F[0] := SRC1.qword[0] G[0] := SRCDST.qword[1] H[0] := SRCDST.qword[0] WK[0] := SRC2.qword[0] WK[1] := SRC2.qword[1] FOR i in 0 to 1: A[i+1] := CH(E[i], F[i], G[i]) + S1(E[i]) + WK[i] + H[i] + MAJ(A[i], B[i], C[i]) + S0(A[i]) B[i+1] := A[i] C[i+1] := B[i] D[i+1] := C[i] E[i+1] := CH(E[i], F[i], G[i]) + S1(E[i]) + WK[i] + H[i] + D[i] F[i+1] := E[i] G[i+1] := F[i] H[i+1] := G[i] SRCDST.qword[3] = A[2] SRCDST.qword[2] = B[2] SRCDST.qword[1] = E[2] SRCDST.qword[0] = F[2]

(143) In certain embodiments, a VSHA512RNDS2 instruction performs (e.g., only) two rounds of SHA512 operation using an initial SHA512 state (C, D, G, H) from the first operand, an initial SHA256 state (A, B, E, F) from the second operand, and a pre-computed sum of the next two round message quadwords and the corresponding round constants (e.g., from the implicit operand xmm0). In one embodiment, only the two lower qwords of XMM0 are used by the instruction. In certain embodiments, an updated SHA512 state (A, B, E, F) is written to the first operand, and the second operand can be used as the updated state (C, D, G, H) in later rounds.

(144) FIG. 19 illustrates circuitry **1900** including an execution circuit **1910** with SHA512 hashing round circuitry **1912** according to embodiments of the disclosure. In certain embodiments, decode circuit (e.g., decode circuit **102** in FIG. 1 or decode circuit **206** in FIG. 2) decodes an instruction into a decoded instruction that causes execution circuit **1910** to perform a hashing operation with SHA512 hashing round circuitry **1912** (for example, the decoded instruction indicates to execution circuit **1910** which components to use, e.g., here to use the SHA512 hashing round circuitry **1912**, the direction(s) of rotation and shifting, and the values of the rotation constants and shifting constants). In the depicted embodiment, an instruction format may include one or more fields that identifies the first packed data destination/source **1901**, second packed data source **1902**, and the third packed data source **1903**. In the depicted embodiment, the first packed data destination/source **1901**, second packed data source **1902**, and the third packed data source **1903** each include four elements (indexed 0-3). However, it should be understood that any number of elements, or a proper subset of any numbers, may be utilized with the instruction(s) disclosed herein. In certain embodiments, each element in a packed data source has a same bit width (e.g., a byte or bytes) and/or each of the elements in each of the other packed data sources.

(145) In one embodiment, first packed data destination/source **1901** is a four element location that stores state elements C, D, G, and H, respectively, second packed data source **1902** is a four element location that stores state elements A, B, E, and F, respectively, and the third packed data source **1903** is a two element location (e.g., the first two elements of a four element location) that stores input message elements (e.g., each element being a pre-computed sum of message elements and corresponding round constants for the next round “0” and the round after next “1”, respectively) (e.g., WK[0], and WK[1]), respectively where “WK” refers to W XOR K for a particular round [round]. In one embodiment, WK[i] uses message word [i] (referred to as “K”) XORed with Key Const[i] (referred to as “K”) for a particular round [i]. In certain embodiments, a message word is source either from (1) the input to a hash (e.g., for the first 16 words) or (2) resultant calculation (e.g., using the VSHA512MSG1/VSHA512MSG2 instructions) for the other the message words (e.g., for the rest 64 words). Input state elements (e.g., C, D, G, and H) may be sourced from a scheduling instruction(s).

(146) In certain embodiments, SHA512 hashing round circuitry **1914** performs hashing operations (e.g., two rounds as discussed herein) on the state elements C, D, G, and H from first packed data destination/source **1901**, the state elements A, B, E, and F from second packed data source **1902**, and a pre-computed sum of the next 2 round message elements (e.g., 64-bit qwords) and the corresponding round constants from the third packed data source **1903** to generate a hash two round result (e.g., including updated state elements A, B, E, and F). The packed data hashing result (e.g., updated state elements A, B, E, and F) may then be stored (e.g., committed to be visible by other instructions) into destination **1901**.

(147) FIG. 20 illustrates a method **2000** of processing an intermediate message instruction according to a SHA512 hashing standard according to embodiments of the disclosure. A processor (e.g., or processor core) may perform method **2000**, e.g., in response to receiving a request to execute an instruction from software. Depicted method **2000** includes processing a SHA512 hashing standard instruction by: fetch a single instruction having an opcode that indicates an intermediate calculation for an input message according to a SHA512 hashing standard is to be performed and at least one first field that identifies five (e.g., 64-bit) message elements W[i+0], W[i+1], W[i+2], W[i+3], and W[i+4] for a round “i” (e.g., an immediately previous round) according to the SHA512 hashing standard (e.g., and a location to write a resultant) at **2002**, decode the instruction into a decoded instruction at **2004**, retrieve data associated with the identified source operands at **2006**, (optionally) schedule the decoded instruction for execution at **2008**, execute the decoded instruction to determine four (e.g., 64-bit) intermediate message elements from the five (e.g., 64-bit) message elements W[i+0], W[i+1], W[i+2], W[i+3], and W[i+4] at **2010** (where “i” is the round number or other scheduling index), and commit the executed instruction to store the four

intermediate message elements into a location specified by the single instruction at **2012**.

(148) The opcode may correspond to (e.g., identifies) the vector instruction mnemonic of VSHA512MSG1 (e.g., where (optional usage of) “V” represents a vector operation (e.g., VEX as discussed below), “SHA512” represents the hashing standard being SHA512, and “MSG1” represents an intermediate message “1” used to generate a final message “2”, with the data from that final message used as an input of a VSHA512RNDS2 instruction).

(149) In one embodiment, an instruction has the following format for its fields: VSHA512MSG1 SOURCE/DESTINATION, SOURCE 1// where source/destination is a first packed data source (e.g., register) and destination, and source 1 is the second packed data source (e.g., which may be a register or memory address). In one embodiment, each of the registers used is a (e.g., 256-bit) YMM register (e.g., YMM1, YMM2, etc.).

(150) Pseudocode for the performance of this instruction may be as follows (e.g., where qword is a quadword size of 64-bits for each element, [x] is an element position indexed with x being from 0 to n, where n-1 is the total number of elements, << is a left rotation (that wraps around) for a given number of bits, >> is a right rotation (that wraps around) for a given number of bits, A is a logical (e.g., binary) XOR operator, | is a logical (e.g., binary) OR operator that outputs a high bit (one) if either bit or both bits in a respective bit position is set to high (one), & is a logical (e.g., binary) AND operator, and ? is a conditional operator (e.g., if condition is true ? then value X: otherwise value Y):

```
(151) TABLE-US-00005    // DEFINITIONS #define ROL64(x, n) (((x) << (n)) {circumflex over ( )} ((x) >> (64-(n)))) #define ROR64(x, n) (((x) >> (n)) {circumflex over ( )} ((x) << (64-(n)))) #define SHR64(x, n) ((x) >> (n)) #define s0(w) (ROR64(w,1) {circumflex over ( )} ROR64(w,8) {circumflex over ( )} SHR64(w,7)) // INSTRUCTION FLOW W[4] := SRC1.qword[0] W[3] := SRCDST.qword[3] W[2] := SRCDST.qword[2] W[1] := SRCDST.qword[1] W[0] := SRCDST.qword[0] SRCDST.qword[3] := W[3] + s0(W[4]) SRCDST.qword[2] := W[2] + s0(W[3]) SRCDST.qword[1] := W[1] + s0(W[2]) SRCDST.qword[0] := W[0] + s0(W[1])
```

(152) In certain embodiments, a VSHA512MSG1 is one of two SHA512 message scheduling instructions. In certain embodiments, the VSHA512MSG1 instruction performs an intermediate calculation for the next four SHA512 message quadwords.

(153) FIG. 21 illustrates circuitry **2100** including an execution circuit **2110** with intermediate SHA512 message calculation circuitry **2112** according to embodiments of the disclosure. In certain embodiments, decode circuit (e.g., decode circuit **102** in FIG. 1 or decode circuit **206** in FIG. 2) decodes an instruction into a decoded instruction that causes execution circuit **2110** to perform a SHA512 message calculation operation with intermediate SHA512 message calculation circuitry **2112** (for example, the decoded instruction indicate to execution circuit **2110** which components to use, e.g., here to use the intermediate SHA512 message calculation circuitry **2112**). In the depicted embodiment, an instruction format may include one or more fields that identifies the first packed data destination/source **2101** and second packed data source **2102**. In the depicted embodiment, the first packed data destination/source **2101** and second packed data source **2102** each include four elements (indexed 0-3). However, it should be understood that any number of elements, or a proper subset of any numbers, may be utilized with the instruction(s) disclosed herein. In certain embodiments, each element in a packed data source has a same bit width (e.g., a byte or bytes) and/or each of the elements in each of the other packed data sources.

(154) In one embodiment, first packed data destination/source **2101** is a four element location that stores input message elements W[i], W[i+1], W[2], and W[i+3], respectively and second packed data source **2102** is a one (e.g., or four elements, with three elements not utilized) element location that stores input message element W[i+4]. Input message elements may be sourced from an input message, e.g., a hashing input. The input message may be from a previous message (e.g., a previous round to the round the current message is being determined).

(155) In certain embodiments, intermediate SHA512 message calculation circuitry **2112** of

execution circuit **2110** performs message generation operations on the input message elements to generate a plurality (e.g., four) intermediate message elements as indicated by the instruction (e.g., by the opcode). In one embodiment, the message generation operations are performed on each element in parallel. The packed data intermediate message result may then be stored (e.g., committed to be visible by other instructions) into destination **2101**, for example, for use by a next (e.g., in program order) of the pair of message scheduling instructions.

(156) FIG. **22** illustrates a method **2200** of processing a final message instruction according to a SHA512 hashing standard according to embodiments of the disclosure. Final message instruction may be next (e.g., in program order) to an intermediate message instruction that is providing its output as an input to the final message instruction. A processor (e.g., or processor core) may perform method **2200**, e.g., in response to receiving a request to execute an instruction from software. Depicted method **2200** includes processing a SHA512 hashing standard instruction by: fetch a single instruction having an opcode that indicates a final calculation for an input message according to a SHA512 hashing standard is to be performed and at least one first field that identifies six (e.g., 64-bit) message elements $W[i+14]$, $W[i+15]$, and four intermediate message elements from an intermediate calculation for a round according to the SHA512 hashing standard at **2222**, decode the instruction into a decoded instruction at **2204**, retrieve data associated with the identified source operands at **2206**, (optionally) schedule the decoded instruction for execution at **2208**, execute the decoded instruction to determine four (e.g., 64-bit) final message elements from the six (e.g., 64-bit) message elements $W[i+14]$, $W[i+15]$, and four intermediate message elements from an intermediate calculation at **2210**, and commit the executed instruction to store the four final message elements (e.g., $W[i+16]$, $W[i+17]$, $W[i+18]$, and $W[i+19]$) into a location specified by the single instruction at **2212**.

(157) The opcode may correspond to (e.g., identifies) the vector instruction mnemonic of VSHA512MSG2 (e.g., where (optional usage of) “V” represents a vector operation (e.g., VEX as discussed below), “SHA512” represents the hashing standard being SHA512, and “MSG1” represents a final message “2” generated from an intermediate message “1”, with the data from the final message used as an input of a VSHA512RND2 instruction).

(158) In one embodiment, an instruction has the following format for its fields: VSHA512MSG2 SOURCE/DESTINATION, SOURCE 1// where source/destination is a first packed data source (e.g., register) and destination, and source 1 is the second packed data source (e.g., which may be a register or memory address). In one embodiment, each of the registers used is a (e.g., 256-bit) YMM register (e.g., YMM1, YMM2, etc.).

(159) Pseudocode for the performance of this instruction may be as follows (e.g., where qword is a quadword size of 62-bits for each element, $[x]$ is an element position indexed with x being from 0 to n , where $n-1$ is the total number of elements, $\ll$ is a left rotation (that wraps around) for a given number of bits, $\gg$ is a right rotation (that wraps around) for a given number of bits, $\wedge$ is a logical (e.g., binary) XOR operator, $\vee$ is a logical (e.g., binary) OR operator that outputs a high bit (one) if either bit or both bits in a respective bit position is set to high (one), $\&$ is a logical (e.g., binary) AND operator, and $?$ is a conditional operator (e.g., if condition is true ? then value X: otherwise value Y):

```
(160) TABLE-US-00006    // DEFINITIONS #define ROL64(x, n) (((x) << (n)) {circumflex over ( )} ((x) >> (64-(n)))) #define ROR64(x, n) (((x) >> (n)) {circumflex over ( )} ((x) << (64-(n)))) #define SHR64(x, n) ((x) >> (n)) #define s1(w) (ROR64(w,19) {circumflex over ( )} ROR64(w, 61) {circumflex over ( )} SHR64(w, 6)) // INSTRUCTION FLOW W[14] := SRC1.qword[2] W[15] := SRC1.qword[3] W[16] := SRCDST.qword[0] + s1(W[14]) W[17] := SRCDST.qword[1] + s1(W[15]) W[18] := SRCDST.qword[2] + s1(W[16]) W[19] := SRCDST.qword[3] + s1(W[17]) SRCDST.qword[3] := W[19] SRCDST.qword[2] := W[18] SRCDST.qword[1] := W[17] SRCDST.qword[0] := W[16]
```

(161) In certain embodiments, a VSHA512MSG2 instruction is one of two SHA2 message

scheduling instructions. In certain embodiments, a VSHA512MSG2 performs the final calculation for the next four SHA512 message quadwords.

(162) FIG. 23 illustrates circuitry 2300 including an execution circuit 2310 with final SHA512 message calculation circuitry 2312 according to embodiments of the disclosure. In certain embodiments, decode circuit (e.g., decode circuit 102 in FIG. 1 or decode circuit 206 in FIG. 2) decodes an instruction into a decoded instruction that causes execution circuit 2310 to perform a SHA512 message calculation operation with final SHA512 message calculation circuitry 2312 (for example, the decoded instruction indicate to execution circuit 2310 which components to use, e.g., here to use the final SHA512 message calculation circuitry 2312). In the depicted embodiment, an instruction format may include one or more fields that identifies the first packed data destination/source 2301 and the second packed data source 2302. In the depicted embodiment, the first packed data destination/source 2301 and the second packed data source 2302 each include four elements (indexed 0-3). However, it should be understood that any number of elements, or a proper subset of any numbers, may be utilized with the instruction(s) disclosed herein. In certain embodiments, each element in a packed data source has a same bit width (e.g., a byte or bytes) and/or each of the elements in each of the other packed data sources.

(163) In one embodiment, first packed data destination/source 2301 is a four element location that stores intermediate message elements from an intermediate SHA512 message calculation instruction (e.g., VSHA512MSG1), respectively and second packed data source 2302 is a two element (e.g., or four elements, with two elements not utilized) location that stores input message elements $W[i+14]$ and $W[i+15]$, respectively. Input message elements may be sourced from an input message, e.g., a hashing input, and a previous intermediate SHA512 message calculation instruction (e.g., VSHA512MSG1). The input message may be from a previous message (e.g., a previous round to the round the current message is being determined).

(164) In certain embodiments, final SHA512 message calculation circuitry 2312 of execution circuit 2310 performs message generation operations on the input message elements and intermediate message elements to generate a plurality (e.g., four) final message elements as indicated by the instruction (e.g., by the opcode). In one embodiment, the message generation operations are performed on each element in parallel. The packed data final message result (e.g., $W[i+16]$, $W[i+17]$, $W[i+18]$, and $W[i+19]$) may then be stored (e.g., committed to be visible by other instructions) into destination 2301, for example, for use by a next (e.g., in program order) of the pair of message scheduling instructions.

(165) At least some embodiments of the disclosed technologies can be described in view of the following examples:

(166) Example 1. A processor comprising: a decode circuit to decode a single instruction into a decoded single instruction, the single instruction comprising at least one first field that identifies eight 32-bit state elements A, B, C, D, E, F, G, and H for a round according to a SM3 hashing standard and at least one second field that identifies an input message; and an execution circuit to execute the decoded single instruction to: rotate state element C left by 9 bits to form a rotated state element C, rotate state element D left by 9 bits to form a rotated state element D, rotate state element G left by 19 bits to form a rotated state element G, rotate state element H left by 19 bits to form a rotated state element H, perform two rounds according to the SM3 hashing standard on the input message and state element A, state element B, rotated state element C, rotated state element D, state element E, state element F, rotated state element G, and rotated state element H to generate an updated state element A, an updated state element B, an updated state element E, and an updated state element F, and store the updated state element A, the updated state element B, the updated state element E, and the updated state element F into a location specified by the single instruction.

(167) Example 2. The processor of example 1, wherein the at least one first field includes: a first register field that identifies a first register that stores state elements C, D, G, and H, and a second register field that identifies a second register that stores state elements A, B, E, and F.

(168) Example 3. The processor of example 2, wherein the at least one second field includes a third register field that identifies the input message.

(169) Example 4. The processor of example 3, wherein the input message is four 32-bit message elements from a zeroth, first, fourth, and fifth element positions of a sixteen 32-bit message element input message.

(170) Example 5. The processor of example 2, wherein the location specified by the single instruction to store the updated state element A, the updated state element B, the updated state element E, and the updated state element F is the first register.

(171) Example 6. The processor of example 1, wherein: the decode circuit is to decode a second instruction into a decoded second instruction, the second instruction comprising at least one field that identifies four 32-bit initial state elements C, D, G, and H for an initial round of a data block according to the SM3 hashing standard; and the execution circuit is to execute the decoded second instruction to: rotate initial state element C right by 9 bits to form the state element C, rotate initial state element D right by 9 bits to form the state element D, rotate initial state element G right by 19 bits to form the state element G, rotate initial state element H right by 19 bits to form the state element H, and store the state element C, the state element D, the state element D, and the state element H in a location specified by the at least one first field of the single instruction.

(172) Example 7. The processor of example 6, wherein the location is a first register.

(173) Example 8. The processor of example 1, wherein: the decode circuit is to decode a second instruction into a decoded second instruction, the second instruction comprising at least one field that identifies the updated state elements C, D, G, and H generated by the single instruction for a final round of a data block according to the SM3 hashing standard; and the execution circuit is to execute the decoded second instruction to: rotate updated state element C left by 9 bits to form a final state element C of the data block, rotate updated state element D left by 9 bits to form a final state element D of the data block, rotate updated state element G left by 19 bits to form a final state element G of the data block, rotate updated state element H left by 19 bits to form a final state element H of the data block, and store the final state element C, the final state element D, the final state element D, and the final state element H in a location specified by the at least one field of the second instruction.

(174) Example 9. A method comprising: decoding a single instruction into a decoded single instruction with a decode circuit of a processor, the single instruction comprising at least one first field that identifies eight 32-bit state elements A, B, C, D, E, F, G, and H for a round according to a SM3 hashing standard and at least one second field that identifies an input message; and executing the decoded single instruction with an execution circuit of the processor to: rotate state element C left by 9 bits to form a rotated state element C, rotate state element D left by 9 bits to form a rotated state element D, rotate state element G left by 19 bits to form a rotated state element G, rotate state element H left by 19 bits to form a rotated state element H, perform two rounds according to the SM3 hashing standard on the input message and state element A, state element B, rotated state element C, rotated state element D, state element E, state element F, rotated state element G, and rotated state element H to generate an updated state element A, an updated state element B, an updated state element E, and an updated state element F, and store the updated state element A, the updated state element B, the updated state element E, and the updated state element F into a location specified by the single instruction.

(175) Example 10. The method of example 9, wherein the at least one first field includes: a first register field that identifies a first register that stores state elements C, D, G, and H, and a second register field that identifies a second register that stores state elements A, B, E, and F.

(176) Example 11. The method of example 10, wherein the at least one second field includes a third register field that identifies the input message.

(177) Example 12. The method of example 11, wherein the input message is four 32-bit message elements from a zeroth, first, fourth, and fifth element positions of a sixteen 32-bit message

element input message.

(178) Example 13. The method of example 10, wherein the location specified by the single instruction to store the updated state element A, the updated state element B, the updated state element E, and the updated state element F is the first register.

(179) Example 14. The method of example 9, further comprising: decoding a second instruction into a decoded second instruction with the decode circuit, the second instruction comprising at least one field that identifies four 32-bit initial state elements C, D, G, and H for an initial round of a data block according to the SM3 hashing standard; and executing the decoded second instruction with the execution circuit to: rotate initial state element C right by 9 bits to form the state element C, rotate initial state element D right by 9 bits to form the state element D, rotate initial state element G right by 19 bits to form the state element G, rotate initial state element H right by 19 bits to form the state element H, and store the state element C, the state element D, the state element D, and the state element H in a location specified by the at least one first field of the single instruction.

(180) Example 15. The method of example 14, further comprising repeating the decoding and the executing of the single instruction for multiple rounds of the data block according to the SM3 hashing standard.

(181) Example 16. The method of example 15, further comprising: decoding a third instruction into a decoded third instruction with the decode circuit, the second instruction comprising at least one field that identifies the updated state elements C, D, G, and H generated by the single instruction for a final round of the multiple rounds of the data block according to the SM3 hashing standard; and executing the third instruction with the execution circuit to: rotate updated state element C left by 9 bits to form a final state element C of the data block, rotate updated state element D left by 9 bits to form a final state element D of the data block, rotate updated state element G left by 19 bits to form a final state element G of the data block, rotate updated state element H left by 19 bits to form a final state element H of the data block, and store the final state element C, the final state element D, the final state element D, and the final state element H in a location specified by the at least one field of the third instruction.

(182) Example 17. A non-transitory machine readable medium that stores code that when executed by a machine causes the machine to perform a method comprising: decoding a single instruction into a decoded single instruction with a decode circuit of a processor, the single instruction comprising at least one first field that identifies eight 32-bit state elements A, B, C, D, E, F, G, and H for a round according to a SM3 hashing standard and at least one second field that identifies an input message; and executing the decoded single instruction with an execution circuit of the processor to: rotate state element C left by 9 bits to form a rotated state element C, rotate state element D left by 9 bits to form a rotated state element D, rotate state element G left by 19 bits to form a rotated state element G, rotate state element H left by 19 bits to form a rotated state element H, perform two rounds according to the SM3 hashing standard on the input message and state element A, state element B, rotated state element C, rotated state element D, state element E, state element F, rotated state element G, and rotated state element H to generate an updated state element A, an updated state element B, an updated state element E, and an updated state element F, and store the updated state element A, the updated state element B, the updated state element E, and the updated state element F into a location specified by the single instruction.

(183) Example 18. The non-transitory machine readable medium of example 17, wherein the at least one first field includes: a first register field that identifies a first register that stores state elements C, D, G, and H, and a second register field that identifies a second register that stores state elements A, B, E, and F.

(184) Example 19. The non-transitory machine readable medium of example 18, wherein the at least one second field includes a third register field that identifies the input message.

(185) Example 20. The non-transitory machine readable medium of example 19, wherein the input message is four 32-bit message elements from a zeroth, first, fourth, and fifth element positions of

a sixteen 32-bit message element input message.

(186) Example 21. The non-transitory machine readable medium of example 18, wherein the location specified by the single instruction to store the updated state element A, the updated state element B, the updated state element E, and the updated state element F is the first register.

(187) Example 22. The non-transitory machine readable medium of example 17, further comprising: decoding a second instruction into a decoded second instruction with the decode circuit, the second instruction comprising at least one field that identifies four 32-bit initial state elements C, D, G, and H for an initial round of a data block according to the SM3 hashing standard; and executing the decoded second instruction with the execution circuit to: rotate initial state element C right by 9 bits to form the state element C, rotate initial state element D right by 9 bits to form the state element D, rotate initial state element G right by 19 bits to form the state element G, rotate initial state element H right by 19 bits to form the state element H, and store the state element C, the state element D, the state element D, and the state element H in a location specified by the at least one first field of the single instruction.

(188) Example 23. The non-transitory machine readable medium of example 22, further comprising repeating the decoding and the executing of the single instruction for multiple rounds of the data block according to the SM3 hashing standard.

(189) Example 24. The non-transitory machine readable medium of example 23, further comprising: decoding a third instruction into a decoded third instruction with the decode circuit, the second instruction comprising at least one field that identifies the updated state elements C, D, G, and H generated by the single instruction for a final round of the multiple rounds of the data block according to the SM3 hashing standard; and executing the third instruction with the execution circuit to: rotate updated state element C left by 9 bits to form a final state element C of the data block, rotate updated state element D left by 9 bits to form a final state element D of the data block, rotate updated state element G left by 19 bits to form a final state element G of the data block, rotate updated state element H left by 19 bits to form a final state element H of the data block, and store the final state element C, the final state element D, the final state element D, and the final state element H in a location specified by the at least one field of the third instruction.

(190) In yet another embodiment, an apparatus comprises a data storage device that stores code that when executed by a hardware processor causes the hardware processor to perform any method disclosed herein. An apparatus may be as described in the detailed description. A method may be as described in the detailed description.

(191) An instruction set may include one or more instruction formats. A given instruction format may define various fields (e.g., number of bits, location of bits) to specify, among other things, the operation to be performed (e.g., opcode) and the operand(s) on which that operation is to be performed and/or other data field(s) (e.g., mask). Some instruction formats are further broken down though the definition of instruction templates (or subformats). For example, the instruction templates of a given instruction format may be defined to have different subsets of the instruction format's fields (the included fields are typically in the same order, but at least some have different bit positions because there are less fields included) and/or defined to have a given field interpreted differently. Thus, each instruction of an ISA is expressed using a given instruction format (and, if defined, in a given one of the instruction templates of that instruction format) and includes fields for specifying the operation and the operands. For example, an exemplary ADD instruction has a specific opcode and an instruction format that includes an opcode field to specify that opcode and operand fields to select operands (source1/destination and source2); and an occurrence of this ADD instruction in an instruction stream will have specific contents in the operand fields that select specific operands. A set of SIMD extensions referred to as the Advanced Vector Extensions (AVX) (AVX1 and AVX2) and using the Vector Extensions (VEX) coding scheme has been released and/or published (e.g., see Intel® 64 and IA-32 Architectures Software Developer's Manual, November 2018; and see Intel® Architecture Instruction Set Extensions Programming Reference,

October 2018).

(192) Exemplary Instruction Formats

(193) Embodiments of the instruction(s) described herein may be embodied in different formats. Additionally, exemplary systems, architectures, and pipelines are detailed below. Embodiments of the instruction(s) may be executed on such systems, architectures, and pipelines, but are not limited to those detailed.

(194) Generic Vector Friendly Instruction Format

(195) A vector friendly instruction format is an instruction format that is suited for vector instructions (e.g., there are certain fields specific to vector operations). While embodiments are described in which both vector and scalar operations are supported through the vector friendly instruction format, alternative embodiments use only vector operations the vector friendly instruction format.

(196) FIGS. **24A-24B** are block diagrams illustrating a generic vector friendly instruction format and instruction templates thereof according to embodiments of the disclosure. FIG. **24A** is a block diagram illustrating a generic vector friendly instruction format and class A instruction templates thereof according to embodiments of the disclosure; while FIG. **24B** is a block diagram illustrating the generic vector friendly instruction format and class B instruction templates thereof according to embodiments of the disclosure. Specifically, a generic vector friendly instruction format **2400** for which are defined class A and class B instruction templates, both of which include no memory access **2405** instruction templates and memory access **2420** instruction templates. The term generic in the context of the vector friendly instruction format refers to the instruction format not being tied to any specific instruction set.

(197) While embodiments of the disclosure will be described in which the vector friendly instruction format supports the following: a 64 byte vector operand length (or size) with 32 bit (4 byte) or 64 bit (8 byte) data element widths (or sizes) (and thus, a 64 byte vector consists of either 16 doubleword-size elements or alternatively, 8 quadword-size elements); a 64 byte vector operand length (or size) with 16 bit (2 byte) or 8 bit (1 byte) data element widths (or sizes); a 32 byte vector operand length (or size) with 32 bit (4 byte), 64 bit (8 byte), 16 bit (2 byte), or 8 bit (1 byte) data element widths (or sizes); and a 16 byte vector operand length (or size) with 32 bit (4 byte), 64 bit (8 byte), 16 bit (2 byte), or 8 bit (1 byte) data element widths (or sizes); alternative embodiments may support more, less and/or different vector operand sizes (e.g., 256 byte vector operands) with more, less, or different data element widths (e.g., 128 bit (16 byte) data element widths).

(198) The class A instruction templates in FIG. **24A** include: 1) within the no memory access **2405** instruction templates there is shown a no memory access, full round control type operation **2410** instruction template and a no memory access, data transform type operation **2415** instruction template; and 2) within the memory access **2420** instruction templates there is shown a memory access, temporal **2425** instruction template and a memory access, non-temporal **2430** instruction template. The class B instruction templates in FIG. **24B** include: 1) within the no memory access **2405** instruction templates there is shown a no memory access, write mask control, partial round control type operation **2412** instruction template and a no memory access, write mask control, vsize type operation **2417** instruction template; and 2) within the memory access **2420** instruction templates there is shown a memory access, write mask control **2427** instruction template.

(199) The generic vector friendly instruction format **2400** includes the following fields listed below in the order illustrated in FIGS. **24A-24B**.

(200) Format field **2440**—a specific value (an instruction format identifier value) in this field uniquely identifies the vector friendly instruction format, and thus occurrences of instructions in the vector friendly instruction format in instruction streams. As such, this field is optional in the sense that it is not needed for an instruction set that has only the generic vector friendly instruction format.

(201) Base operation field **2442**—its content distinguishes different base operations.

(202) Register index field **2444**—its content, directly or through address generation, specifies the locations of the source and destination operands, be they in registers or in memory. These include a sufficient number of bits to select N registers from a $P \times Q$ (e.g. 32×512 , 16×128 , 32×1024 , 64×1024) register file. While in one embodiment N may be up to three sources and one destination register, alternative embodiments may support more or less sources and destination registers (e.g., may support up to two sources where one of these sources also acts as the destination, may support up to three sources where one of these sources also acts as the destination, may support up to two sources and one destination).

(203) Modifier field **2446**—its content distinguishes occurrences of instructions in the generic vector instruction format that specify memory access from those that do not; that is, between no memory access **2405** instruction templates and memory access **2420** instruction templates. Memory access operations read and/or write to the memory hierarchy (in some cases specifying the source and/or destination addresses using values in registers), while non-memory access operations do not (e.g., the source and destinations are registers). While in one embodiment this field also selects between three different ways to perform memory address calculations, alternative embodiments may support more, less, or different ways to perform memory address calculations.

(204) Augmentation operation field **2450**—its content distinguishes which one of a variety of different operations to be performed in addition to the base operation. This field is context specific. In one embodiment of the disclosure, this field is divided into a class field **2468**, an alpha field **2452**, and a beta field **2454**. The augmentation operation field **2450** allows common groups of operations to be performed in a single instruction rather than 2, 3, or 4 instructions.

(205) Scale field **2460**—its content allows for the scaling of the index field's content for memory address generation (e.g., for address generation that uses $2.\text{sup.scale} \times \text{index} + \text{base}$).

(206) Displacement Field **2462A**—its content is used as part of memory address generation (e.g., for address generation that uses $2.\text{sup.scale} \times \text{index} + \text{base} + \text{displacement}$).

(207) Displacement Factor Field **2462B** (note that the juxtaposition of displacement field **2462A** directly over displacement factor field **2462B** indicates one or the other is used) —its content is used as part of address generation; it specifies a displacement factor that is to be scaled by the size of a memory access (N) —where N is the number of bytes in the memory access (e.g., for address generation that uses $2.\text{sup.scale} \times \text{index} + \text{base} + \text{scaled displacement}$). Redundant low-order bits are ignored and hence, the displacement factor field's content is multiplied by the memory operands total size (N) in order to generate the final displacement to be used in calculating an effective address. The value of N is determined by the processor hardware at runtime based on the full opcode field **2474** (described later herein) and the data manipulation field **2454C**. The displacement field **2462A** and the displacement factor field **2462B** are optional in the sense that they are not used for the no memory access **2405** instruction templates and/or different embodiments may implement only one or none of the two.

(208) Data element width field **2464**—its content distinguishes which one of a number of data element widths is to be used (in some embodiments for all instructions; in other embodiments for only some of the instructions). This field is optional in the sense that it is not needed if only one data element width is supported and/or data element widths are supported using some aspect of the opcodes.

(209) Write mask field **2470**—its content controls, on a per data element position basis, whether that data element position in the destination vector operand reflects the result of the base operation and augmentation operation. Class A instruction templates support merging-writemasking, while class B instruction templates support both merging- and zeroing-writemasking. When merging, vector masks allow any set of elements in the destination to be protected from updates during the execution of any operation (specified by the base operation and the augmentation operation); in other one embodiment, preserving the old value of each element of the destination where the corresponding mask bit has a 0. In contrast, when zeroing vector masks allow any set of elements

in the destination to be zeroed during the execution of any operation (specified by the base operation and the augmentation operation); in one embodiment, an element of the destination is set to 0 when the corresponding mask bit has a 0 value. A subset of this functionality is the ability to control the vector length of the operation being performed (that is, the span of elements being modified, from the first to the last one); however, it is not necessary that the elements that are modified be consecutive. Thus, the write mask field **2470** allows for partial vector operations, including loads, stores, arithmetic, logical, etc. While embodiments of the disclosure are described in which the write mask field's **2470** content selects one of a number of write mask registers that contains the write mask to be used (and thus the write mask field's **2470** content indirectly identifies that masking to be performed), alternative embodiments instead or additionally allow the mask write field's **2470** content to directly specify the masking to be performed.

(210) Immediate field **2472**—its content allows for the specification of an immediate. This field is optional in the sense that it is not present in an implementation of the generic vector friendly format that does not support immediate and it is not present in instructions that do not use an immediate.

(211) Class field **2468**—its content distinguishes between different classes of instructions. With reference to FIGS. **24A-B**, the contents of this field select between class A and class B instructions. In FIGS. **24A-B**, rounded corner squares are used to indicate a specific value is present in a field (e.g., class A **2468A** and class B **2468B** for the class field **2468** respectively in FIGS. **24A-B**).

(212) Instruction Templates of Class A

(213) In the case of the non-memory access **2405** instruction templates of class A, the alpha field **2452** is interpreted as an RS field **2452A**, whose content distinguishes which one of the different augmentation operation types are to be performed (e.g., round **2452A.1** and data transform **2452A.2** are respectively specified for the no memory access, round type operation **2410** and the no memory access, data transform type operation **2415** instruction templates), while the beta field **2454** distinguishes which of the operations of the specified type is to be performed. In the no memory access **2405** instruction templates, the scale field **2460**, the displacement field **2462A**, and the displacement scale field **2462B** are not present.

(214) No-Memory Access Instruction Templates—Full Round Control Type Operation

(215) In the no memory access full round control type operation **2410** instruction template, the beta field **2454** is interpreted as a round control field **2454A**, whose content(s) provide static rounding. While in the described embodiments of the disclosure the round control field **2454A** includes a suppress all floating point exceptions (SAE) field **2456** and a round operation control field **2458**, alternative embodiments may support may encode both these concepts into the same field or only have one or the other of these concepts/fields (e.g., may have only the round operation control field **2458**).

(216) SAE field **2456**—its content distinguishes whether or not to disable the exception event reporting; when the SAE field's **2456** content indicates suppression is enabled, a given instruction does not report any kind of floating-point exception flag and does not raise any floating point exception handler.

(217) Round operation control field **2458**—its content distinguishes which one of a group of rounding operations to perform (e.g., Round-up, Round-down, Round-towards-zero and Round-to-nearest). Thus, the round operation control field **2458** allows for the changing of the rounding mode on a per instruction basis. In one embodiment of the disclosure where a processor includes a control register for specifying rounding modes, the round operation control field's **2450** content overrides that register value.

(218) No Memory Access Instruction Templates—Data Transform Type Operation

(219) In the no memory access data transform type operation **2415** instruction template, the beta field **2454** is interpreted as a data transform field **2454B**, whose content distinguishes which one of a number of data transforms is to be performed (e.g., no data transform, swizzle, broadcast).

(220) In the case of a memory access **2420** instruction template of class A, the alpha field **2452** is

interpreted as an eviction hint field **2452B**, whose content distinguishes which one of the eviction hints is to be used (in FIG. **24A**, temporal **2452B.1** and non-temporal **2452B.2** are respectively specified for the memory access, temporal **2425** instruction template and the memory access, non-temporal **2430** instruction template), while the beta field **2454** is interpreted as a data manipulation field **2454C**, whose content distinguishes which one of a number of data manipulation operations (also known as primitives) is to be performed (e.g., no manipulation; broadcast; up conversion of a source; and down conversion of a destination). The memory access **2420** instruction templates include the scale field **2460**, and optionally the displacement field **2462A** or the displacement scale field **2462B**.

(221) Vector memory instructions perform vector loads from and vector stores to memory, with conversion support. As with regular vector instructions, vector memory instructions transfer data from/to memory in a data element-wise fashion, with the elements that are actually transferred being dictated by the contents of the vector mask that is selected as the write mask.

(222) Memory Access Instruction Templates—Temporal

(223) Temporal data is data likely to be reused soon enough to benefit from caching. This is, however, a hint, and different processors may implement it in different ways, including ignoring the hint entirely.

(224) Memory Access Instruction Templates—Non-Temporal

(225) Non-temporal data is data unlikely to be reused soon enough to benefit from caching in the 1st-level cache and should be given priority for eviction. This is, however, a hint, and different processors may implement it in different ways, including ignoring the hint entirely.

(226) Instruction Templates of Class B

(227) In the case of the instruction templates of class B, the alpha field **2452** is interpreted as a write mask control (Z) field **2452C**, whose content distinguishes whether the write masking controlled by the write mask field **2470** should be a merging or a zeroing.

(228) In the case of the non-memory access **2405** instruction templates of class B, part of the beta field **2454** is interpreted as an RL field **2457A**, whose content distinguishes which one of the different augmentation operation types are to be performed (e.g., round **2457A.1** and vector length (VSIZE) **2457A.2** are respectively specified for the no memory access, write mask control, partial round control type operation **2412** instruction template and the no memory access, write mask control, VSIZE type operation **2417** instruction template), while the rest of the beta field **2454** distinguishes which of the operations of the specified type is to be performed. In the no memory access **2405** instruction templates, the scale field **2460**, the displacement field **2462A**, and the displacement scale field **2462B** are not present.

(229) In the no memory access, write mask control, partial round control type operation **2410** instruction template, the rest of the beta field **2454** is interpreted as a round operation field **2459A** and exception event reporting is disabled (a given instruction does not report any kind of floating-point exception flag and does not raise any floating point exception handler).

(230) Round operation control field **2459A**—just as round operation control field **2458**, its content distinguishes which one of a group of rounding operations to perform (e.g., Round-up, Round-down, Round-towards-zero and Round-to-nearest). Thus, the round operation control field **2459A** allows for the changing of the rounding mode on a per instruction basis. In one embodiment of the disclosure where a processor includes a control register for specifying rounding modes, the round operation control field's **2450** content overrides that register value.

(231) In the no memory access, write mask control, VSIZE type operation **2417** instruction template, the rest of the beta field **2454** is interpreted as a vector length field **2459B**, whose content distinguishes which one of a number of data vector lengths is to be performed on (e.g., 128, 256, or 512 byte).

(232) In the case of a memory access **2420** instruction template of class B, part of the beta field **2454** is interpreted as a broadcast field **2457B**, whose content distinguishes whether or not the

broadcast type data manipulation operation is to be performed, while the rest of the beta field **2454** is interpreted the vector length field **2459B**. The memory access **2420** instruction templates include the scale field **2460**, and optionally the displacement field **2462A** or the displacement scale field **2462B**.

(233) With regard to the generic vector friendly instruction format **2400**, a full opcode field **2474** is shown including the format field **2440**, the base operation field **2442**, and the data element width field **2464**. While one embodiment is shown where the full opcode field **2474** includes all of these fields, the full opcode field **2474** includes less than all of these fields in embodiments that do not support all of them. The full opcode field **2474** provides the operation code (opcode).

(234) The augmentation operation field **2450**, the data element width field **2464**, and the write mask field **2470** allow these features to be specified on a per instruction basis in the generic vector friendly instruction format.

(235) The combination of write mask field and data element width field create typed instructions in that they allow the mask to be applied based on different data element widths.

(236) The various instruction templates found within class A and class B are beneficial in different situations. In some embodiments of the disclosure, different processors or different cores within a processor may support only class A, only class B, or both classes. For instance, a high performance general purpose out-of-order core intended for general-purpose computing may support only class B, a core intended primarily for graphics and/or scientific (throughput) computing may support only class A, and a core intended for both may support both (of course, a core that has some mix of templates and instructions from both classes but not all templates and instructions from both classes is within the purview of the disclosure). Also, a single processor may include multiple cores, all of which support the same class or in which different cores support different class. For instance, in a processor with separate graphics and general purpose cores, one of the graphics cores intended primarily for graphics and/or scientific computing may support only class A, while one or more of the general purpose cores may be high performance general purpose cores with out of order execution and register renaming intended for general-purpose computing that support only class B. Another processor that does not have a separate graphics core, may include one more general purpose in-order or out-of-order cores that support both class A and class B. Of course, features from one class may also be implemented in the other class in different embodiments of the disclosure. Programs written in a high level language would be put (e.g., just in time compiled or statically compiled) into an variety of different executable forms, including: 1) a form having only instructions of the class(es) supported by the target processor for execution; or 2) a form having alternative routines written using different combinations of the instructions of all classes and having control flow code that selects the routines to execute based on the instructions supported by the processor which is currently executing the code.

(237) Exemplary Specific Vector Friendly Instruction Format

(238) FIG. **25** is a block diagram illustrating an exemplary specific vector friendly instruction format according to embodiments of the disclosure. FIG. **25** shows a specific vector friendly instruction format **2500** that is specific in the sense that it specifies the location, size, interpretation, and order of the fields, as well as values for some of those fields. The specific vector friendly instruction format **2500** may be used to extend the x86 instruction set, and thus some of the fields are similar or the same as those used in the existing x86 instruction set and extension thereof (e.g., AVX). This format remains consistent with the prefix encoding field, real opcode byte field, MOD R/M field, SIB field, displacement field, and immediate fields of the existing x86 instruction set with extensions. The fields from FIG. **24** into which the fields from FIG. **25** map are illustrated.

(239) It should be understood that, although embodiments of the disclosure are described with reference to the specific vector friendly instruction format **2500** in the context of the generic vector friendly instruction format **2400** for illustrative purposes, the disclosure is not limited to the specific vector friendly instruction format **2500** except where claimed. For example, the generic

vector friendly instruction format **2400** contemplates a variety of possible sizes for the various fields, while the specific vector friendly instruction format **2500** is shown as having fields of specific sizes. By way of specific example, while the data element width field **2464** is illustrated as a one bit field in the specific vector friendly instruction format **2500**, the disclosure is not so limited (that is, the generic vector friendly instruction format **2400** contemplates other sizes of the data element width field **2464**).

(240) The generic vector friendly instruction format **2400** includes the following fields listed below in the order illustrated in FIG. 25A.

(241) EVEX Prefix (Bytes 0-3) **2502**—is encoded in a four-byte form.

(242) Format Field **2440** (EVEX Byte 0, bits [7:0]) —the first byte (EVEX Byte 0) is the format field **2440** and it contains 0x62 (the unique value used for distinguishing the vector friendly instruction format in one embodiment of the disclosure).

(243) The second-fourth bytes (EVEX Bytes 1-3) include a number of bit fields providing specific capability.

(244) REX field **2505** (EVEX Byte 1, bits [7-5]) —consists of an EVEX.R bit field (EVEX Byte 1, bit [7]-R), EVEX.X bit field (EVEX byte 1, bit [6]-X), and 2457BEX byte 1, bit[5]-B). The EVEX.R, EVEX.X, and EVEX.B bit fields provide the same functionality as the corresponding VEX bit fields, and are encoded using is complement form, i.e. ZMM0 is encoded as 1111B, ZMM15 is encoded as 0000B. Other fields of the instructions encode the lower three bits of the register indexes as is known in the art (rrr, xxx, and bbb), so that Rrrr, Xxxx, and Bbbb may be formed by adding EVEX.R, EVEX.X, and EVEX.B.

(245) REX' field **2410**—this is the first part of the REX' field **2410** and is the EVEX.R' bit field (EVEX Byte 1, bit [4]-R') that is used to encode either the upper 16 or lower 16 of the extended 32 register set. In one embodiment of the disclosure, this bit, along with others as indicated below, is stored in bit inverted format to distinguish (in the well-known x86 32-bit mode) from the BOUND instruction, whose real opcode byte is 62, but does not accept in the MOD R/M field (described below) the value of 11 in the MOD field; alternative embodiments of the disclosure do not store this and the other indicated bits below in the inverted format. A value of 1 is used to encode the lower 16 registers. In other words, R'Rrrr is formed by combining EVEX.R', EVEX.R, and the other RRR from other fields.

(246) Opcode map field **2515** (EVEX byte 1, bits [3:0]-mmmm) —its content encodes an implied leading opcode byte (0F, 0F 38, or 0F 3).

(247) Data element width field **2464** (EVEX byte 2, bit [7]-W) —is represented by the notation EVEX.W. EVEX.W is used to define the granularity (size) of the datatype (either 32-bit data elements or 64-bit data elements).

(248) EVEX.vvvv **2520** (EVEX Byte 2, bits [6:3]-vvvv) —the role of EVEX.vvvv may include the following: 1) EVEX.vvvv encodes the first source register operand, specified in inverted (1s complement) form and is valid for instructions with 2 or more source operands; 2) EVEX.vvvv encodes the destination register operand, specified in is complement form for certain vector shifts; or 3) EVEX.vvvv does not encode any operand, the field is reserved and should contain 1111b. Thus, EVEX.vvvv field **2520** encodes the 4 low-order bits of the first source register specifier stored in inverted (1s complement) form. Depending on the instruction, an extra different EVEX bit field is used to extend the specifier size to 32 registers.

(249) EVEX.U 2468 Class field (EVEX byte 2, bit [2]-U) —If EVEX.U=0, it indicates class A or EVEX.U0; if EVEX.U=1, it indicates class B or EVEX.U1.

(250) Prefix encoding field **2525** (EVEX byte 2, bits [1:0]-pp) —provides additional bits for the base operation field. In addition to providing support for the legacy SSE instructions in the EVEX prefix format, this also has the benefit of compacting the SIMD prefix (rather than requiring a byte to express the SIMD prefix, the EVEX prefix requires only 2 bits). In one embodiment, to support legacy SSE instructions that use a SIMD prefix (66H, F2H, F3H) in both the legacy format and in

the EVEX prefix format, these legacy SIMD prefixes are encoded into the SIMD prefix encoding field; and at runtime are expanded into the legacy SIMD prefix prior to being provided to the decode circuit's PLA (so the PLA can execute both the legacy and EVEX format of these legacy instructions without modification). Although newer instructions could use the EVEX prefix encoding field's content directly as an opcode extension, certain embodiments expand in a similar fashion for consistency but allow for different meanings to be specified by these legacy SIMD prefixes. An alternative embodiment may redesign the PLA to support the 2 bit SIMD prefix encodings, and thus not require the expansion.

(251) Alpha field **2452** (EVEX byte 3, bit [7]-EH; also known as EVEX.EH, EVEX.rs, EVEX.RL, EVEX.write mask control, and EVEX.N; also illustrated with α) —as previously described, this field is context specific.

(252) Beta field **2454** (EVEX byte 3, bits [6:4]-SSS, also known as EVEX.s.sub.2-0, EVEX.r.sub.2-0, EVEX.rr1, EVEX.LL0, EVEX.LLB; also illustrated with $\beta\beta\beta$) —as previously described, this field is context specific.

(253) REX' field **2410**—this is the remainder of the REX' field and is the EVEX.V' bit field (EVEX Byte 3, bit [3]-V') that may be used to encode either the upper 16 or lower 16 of the extended 32 register set. This bit is stored in bit inverted format. A value of 1 is used to encode the lower 16 registers. In other words, V'VVVV is formed by combining EVEX.V', EVEX.vvvv.

(254) Write mask field **2470** (EVEX byte 3, bits [2:0]-kkk) —its content specifies the index of a register in the write mask registers as previously described. In one embodiment of the disclosure, the specific value EVEX kkk=000 has a special behavior implying no write mask is used for the particular instruction (this may be implemented in a variety of ways including the use of a write mask hardwired to all ones or hardware that bypasses the masking hardware).

(255) Real Opcode Field **2530** (Byte 4) is also known as the opcode byte. Part of the opcode is specified in this field.

(256) MOD R/M Field **2540** (Byte 5) includes MOD field **2542**, Reg field **2544**, and R/M field **2546**. As previously described, the MOD field's **2542** content distinguishes between memory access and non-memory access operations. The role of Reg field **2544** can be summarized to two situations: encoding either the destination register operand or a source register operand, or be treated as an opcode extension and not used to encode any instruction operand. The role of R/M field **2546** may include the following: encoding the instruction operand that references a memory address, or encoding either the destination register operand or a source register operand.

(257) Scale, Index, Base (SIB) Byte (Byte 6) —As previously described, the scale field's **2450** content is used for memory address generation. SIB.xxx **2554** and SIB.bbb **2556**—the contents of these fields have been previously referred to with regard to the register indexes Xxxx and Bbbb.

(258) Displacement field **2462A** (Bytes 7-10) —when MOD field **2542** contains 10, bytes 7-10 are the displacement field **2462A**, and it works the same as the legacy 32-bit displacement (disp32) and works at byte granularity.

(259) Displacement factor field **2462B** (Byte 7) —when MOD field **2542** contains 01, byte 7 is the displacement factor field **2462B**. The location of this field is that same as that of the legacy x86 instruction set 8-bit displacement (disp8), which works at byte granularity. Since disp8 is sign extended, it can only address between -128 and 127 bytes offsets; in terms of 64 byte cache lines, disp8 uses 8 bits that can be set to only four really useful values -128, -64, 0, and 64; since a greater range is often needed, disp32 is used; however, disp32 requires 4 bytes. In contrast to disp8 and disp32, the displacement factor field **2462B** is a reinterpretation of disp8; when using displacement factor field **2462B**, the actual displacement is determined by the content of the displacement factor field multiplied by the size of the memory operand access (N). This type of displacement is referred to as disp8*N. This reduces the average instruction length (a single byte of used for the displacement but with a much greater range). Such compressed displacement is based on the assumption that the effective displacement is multiple of the granularity of the memory

access, and hence, the redundant low-order bits of the address offset do not need to be encoded. In other words, the displacement factor field **2462B** substitutes the legacy x86 instruction set 8-bit displacement. Thus, the displacement factor field **2462B** is encoded the same way as an x86 instruction set 8-bit displacement (so no changes in the ModRM/SIB encoding rules) with the only exception that disp8 is overloaded to disp8*N. In other words, there are no changes in the encoding rules or encoding lengths but only in the interpretation of the displacement value by hardware (which needs to scale the displacement by the size of the memory operand to obtain a byte-wise address offset). Immediate field **2472** operates as previously described.

(260) Full Opcode Field

(261) FIG. **25B** is a block diagram illustrating the fields of the specific vector friendly instruction format **2500** that make up the full opcode field **2474** according to one embodiment of the disclosure. Specifically, the full opcode field **2474** includes the format field **2440**, the base operation field **2442**, and the data element width (W) field **2464**. The base operation field **2442** includes the prefix encoding field **2525**, the opcode map field **2515**, and the real opcode field **2530**.

(262) Register Index Field

(263) FIG. **25C** is a block diagram illustrating the fields of the specific vector friendly instruction format **2500** that make up the register index field **2444** according to one embodiment of the disclosure. Specifically, the register index field **2444** includes the REX field **2505**, the REX' field **2510**, the MODR/M.reg field **2544**, the MODR/M.r/m field **2546**, the VVVV field **2520**, xxx field **2554**, and the bbb field **2556**.

(264) Augmentation Operation Field

(265) FIG. **25D** is a block diagram illustrating the fields of the specific vector friendly instruction format **2500** that make up the augmentation operation field **2450** according to one embodiment of the disclosure. When the class (U) field **2468** contains 0, it signifies EVEX.U0 (class A **2468A**); when it contains 1, it signifies EVEX.U1 (class B **2468B**). When U=0 and the MOD field **2542** contains 11 (signifying a no memory access operation), the alpha field **2452** (EVEX byte 3, bit [7]-EH) is interpreted as the rs field **2452A**. When the rs field **2452A** contains a 1 (round **2452A.1**), the beta field **2454** (EVEX byte 3, bits [6:4]-SSS) is interpreted as the round control field **2454A**. The round control field **2454A** includes a one bit SAE field **2456** and a two bit round operation field **2458**. When the rs field **2452A** contains a 0 (data transform **2452A.2**), the beta field **2454** (EVEX byte 3, bits [6:4]-SSS) is interpreted as a three bit data transform field **2454B**. When U=0 and the MOD field **2542** contains 00, 01, or 10 (signifying a memory access operation), the alpha field **2452** (EVEX byte 3, bit [7]-EH) is interpreted as the eviction hint (EH) field **2452B** and the beta field **2454** (EVEX byte 3, bits [6:4]-SSS) is interpreted as a three bit data manipulation field **2454C**.

(266) When U=1, the alpha field **2452** (EVEX byte 3, bit [7]-EH) is interpreted as the write mask control (Z) field **2452C**. When U=1 and the MOD field **2542** contains 11 (signifying a no memory access operation), part of the beta field **2454** (EVEX byte 3, bit [4]-S.sub.0) is interpreted as the RL field **2457A**; when it contains a 1 (round **2457A.1**) the rest of the beta field **2454** (EVEX byte 3, bit [6-5]-S.sub.2-1) is interpreted as the round operation field **2459A**, while when the RL field **2457A** contains a 0 (VSIZE **2457A.2**) the rest of the beta field **2454** (EVEX byte 3, bit [6-5]-S.sub.2-1) is interpreted as the vector length field **2459B** (EVEX byte 3, bit [6-5]-L.sub.1-0). When U=1 and the MOD field **2542** contains 00, 01, or 10 (signifying a memory access operation), the beta field **2454** (EVEX byte 3, bits [6:4]-SSS) is interpreted as the vector length field **2459B** (EVEX byte 3, bit [6-5]-L.sub.1-0) and the broadcast field **2457B** (EVEX byte 3, bit [4]-B).

(267) Exemplary Register Architecture

(268) FIG. **26** is a block diagram of a register architecture **2600** according to one embodiment of the disclosure. In the embodiment illustrated, there are 32 vector registers **2610** that are 512 bits wide; these registers are referenced as zmm0 through zmm31. The lower order 256 bits of the lower 16 zmm registers are overlaid on registers ymm0-16. The lower order 128 bits of the lower

16 zmm registers (the lower order 128 bits of the ymm registers) are overlaid on registers xmm0-15. The specific vector friendly instruction format **2500** operates on these overlaid register file as illustrated in the below tables.

(269) TABLE-US-00007 Adjustable Vector Length Class Operations Registers Instruction

Templates that do A (FIG. 24A; 2410, 2415, zmm registers (the vector not include the vector length U = 0) 2425, 2430 length is 64 byte) field 2459B B (FIG. 24B; 2412 zmm registers (the vector U = 1) length is 64 byte) Instruction templates that do B (FIG. 24B; 2417, 2427 zmm, ymm, or xmm include the vector length field U = 1) registers (the vector length is 2459B 64 byte, 32 byte, or 16 byte) depending on the vector length field 2459B

(270) In other words, the vector length field **2459B** selects between a maximum length and one or more other shorter lengths, where each such shorter length is half the length of the preceding length; and instructions templates without the vector length field **2459B** operate on the maximum vector length. Further, in one embodiment, the class B instruction templates of the specific vector friendly instruction format **2500** operate on packed or scalar single/double-precision floating point data and packed or scalar integer data. Scalar operations are operations performed on the lowest order data element position in a zmm/ymm/xmm register; the higher order data element positions are either left the same as they were prior to the instruction or zeroed depending on the embodiment.

(271) Write mask registers **2615**—in the embodiment illustrated, there are 8 write mask registers (k0 through k7), each 64 bits in size. In an alternate embodiment, the write mask registers **2615** are 16 bits in size. As previously described, in one embodiment of the disclosure, the vector mask register k0 cannot be used as a write mask; when the encoding that would normally indicate k0 is used for a write mask, it selects a hardwired write mask of 0xFFFF, effectively disabling write masking for that instruction.

(272) General-purpose registers **2625**—in the embodiment illustrated, there are sixteen 64-bit general-purpose registers that are used along with the existing x86 addressing modes to address memory operands. These registers are referenced by the names RAX, RBX, RCX, RDX, RBP, RSI, RDI, RSP, and R8 through R15.

(273) Scalar floating point stack register file (x87 stack) **2645**, on which is aliased the MMX packed integer flat register file **2650**—in the embodiment illustrated, the x87 stack is an eight-element stack used to perform scalar floating-point operations on 32/64/80-bit floating point data using the x87 instruction set extension; while the MMX registers are used to perform operations on 64-bit packed integer data, as well as to hold operands for some operations performed between the MMX and XMM registers.

(274) Alternative embodiments of the disclosure may use wider or narrower registers. Additionally, alternative embodiments of the disclosure may use more, less, or different register files and registers.

(275) Exemplary Core Architectures, Processors, and Computer Architectures

(276) Processor cores may be implemented in different ways, for different purposes, and in different processors. For instance, implementations of such cores may include: 1) a general purpose in-order core intended for general-purpose computing; 2) a high performance general purpose out-of-order core intended for general-purpose computing; 3) a special purpose core intended primarily for graphics and/or scientific (throughput) computing. Implementations of different processors may include: 1) a CPU including one or more general purpose in-order cores intended for general-purpose computing and/or one or more general purpose out-of-order cores intended for general-purpose computing; and 2) a coprocessor including one or more special purpose cores intended primarily for graphics and/or scientific (throughput). Such different processors lead to different computer system architectures, which may include: 1) the coprocessor on a separate chip from the CPU; 2) the coprocessor on a separate die in the same package as a CPU; 3) the coprocessor on the same die as a CPU (in which case, such a coprocessor is sometimes referred to as special purpose

logic, such as integrated graphics and/or scientific (throughput) logic, or as special purpose cores); and 4) a system on a chip that may include on the same die the described CPU (sometimes referred to as the application core(s) or application processor(s)), the above described coprocessor, and additional functionality. Exemplary core architectures are described next, followed by descriptions of exemplary processors and computer architectures.

(277) Exemplary Core Architectures

(278) In-Order and Out-of-Order Core Block Diagram

(279) FIG. 27A is a block diagram illustrating both an exemplary in-order pipeline and an exemplary register renaming, out-of-order issue/execution pipeline according to embodiments of the disclosure. FIG. 27B is a block diagram illustrating both an exemplary embodiment of an in-order architecture core and an exemplary register renaming, out-of-order issue/execution architecture core to be included in a processor according to embodiments of the disclosure. The solid lined boxes in FIGS. 27A-B illustrate the in-order pipeline and in-order core, while the optional addition of the dashed lined boxes illustrates the register renaming, out-of-order issue/execution pipeline and core. Given that the in-order aspect is a subset of the out-of-order aspect, the out-of-order aspect will be described.

(280) In FIG. 27A, a processor pipeline 2700 includes a fetch stage 2702, a length decode stage 2704, a decode stage 2706, an allocation stage 2708, a renaming stage 2710, a scheduling (also known as a dispatch or issue) stage 2712, a register read/memory read stage 2714, an execute stage 2716, a write back/memory write stage 2718, an exception handling stage 2722, and a commit stage 2724.

(281) FIG. 27B shows processor core 2790 including a front end unit 2730 coupled to an execution engine unit 2750, and both are coupled to a memory unit 2770. The core 2790 may be a reduced instruction set computing (RISC) core, a complex instruction set computing (CISC) core, a very long instruction word (VLIW) core, or a hybrid or alternative core type. As yet another option, the core 2790 may be a special-purpose core, such as, for example, a network or communication core, compression engine, coprocessor core, general purpose computing graphics processing unit (GPGPU) core, graphics core, or the like.

(282) The front end unit 2730 includes a branch prediction unit 2732 coupled to an instruction cache unit 2734, which is coupled to an instruction translation lookaside buffer (TLB) 2736, which is coupled to an instruction fetch unit 2738, which is coupled to a decode unit 2740. The decode unit 2740 (e.g., decode circuit) may decode instructions (e.g., macro-instructions), and generate as an output one or more micro-operations, micro-code entry points, micro-instructions, other instructions, or other control signals, which are decoded from, or which otherwise reflect, or are derived from, the original instructions. The decode unit 2740 may be implemented using various different mechanisms. Examples of suitable mechanisms include, but are not limited to, look-up tables, hardware implementations, programmable logic arrays (PLAs), microcode read only memories (ROMs), etc. In one embodiment, the core 2790 includes a microcode ROM or other medium that stores microcode for certain macro-instructions (e.g., in decode unit 2740 or otherwise within the front end unit 2730). The decode unit 2740 is coupled to a rename/allocator unit 2752 in the execution engine unit 2750.

(283) The execution engine unit 2750 includes the rename/allocator unit 2752 coupled to a retirement unit 2754 and a set of one or more scheduler unit(s) 2756. The scheduler unit(s) 2756 represents any number of different schedulers, including reservations stations, central instruction window, etc. The scheduler unit(s) 2756 is coupled to the physical register file(s) unit(s) 2758. Each of the physical register file(s) units 2758 represents one or more physical register files, different ones of which store one or more different data types, such as scalar integer, scalar floating point, packed integer, packed floating point, vector integer, vector floating point, status (e.g., an instruction pointer that is the address of the next instruction to be executed), etc. In one embodiment, the physical register file(s) unit 2758 comprises a vector registers unit, a write mask

registers unit, and a scalar registers unit. These register units may provide architectural vector registers, vector mask registers, and general purpose registers. The physical register file(s) unit(s) **2758** is overlapped by the retirement unit **2754** to illustrate various ways in which register renaming and out-of-order execution may be implemented (e.g., using a reorder buffer(s) and a retirement register file(s); using a future file(s), a history buffer(s), and a retirement register file(s); using a register maps and a pool of registers; etc.). The retirement unit **2754** and the physical register file(s) unit(s) **2758** are coupled to the execution cluster(s) **2760**. The execution cluster(s) **2760** includes a set of one or more execution units **2762** (e.g., execution circuits) and a set of one or more memory access units **2764**. The execution units **2762** may perform various operations (e.g., shifts, addition, subtraction, multiplication) and on various types of data (e.g., scalar floating point, packed integer, packed floating point, vector integer, vector floating point). While some embodiments may include a number of execution units dedicated to specific functions or sets of functions, other embodiments may include only one execution unit or multiple execution units that all perform all functions. The scheduler unit(s) **2756**, physical register file(s) unit(s) **2758**, and execution cluster(s) **2760** are shown as being possibly plural because certain embodiments create separate pipelines for certain types of data/operations (e.g., a scalar integer pipeline, a scalar floating point/packed integer/packed floating point/vector integer/vector floating point pipeline, and/or a memory access pipeline that each have their own scheduler unit, physical register file(s) unit, and/or execution cluster—and in the case of a separate memory access pipeline, certain embodiments are implemented in which only the execution cluster of this pipeline has the memory access unit(s) **2764**). It should also be understood that where separate pipelines are used, one or more of these pipelines may be out-of-order issue/execution and the rest in-order.

(284) The set of memory access units **2764** is coupled to the memory unit **2770**, which includes a data TLB unit **2772** coupled to a data cache unit **2774** coupled to a level 2 (L2) cache unit **2776**. In one exemplary embodiment, the memory access units **2764** may include a load unit, a store address unit, and a store data unit, each of which is coupled to the data TLB unit **2772** in the memory unit **2770**. The instruction cache unit **2734** is further coupled to a level 2 (L2) cache unit **2776** in the memory unit **2770**. The L2 cache unit **2776** is coupled to one or more other levels of cache and eventually to a main memory.

(285) By way of example, the exemplary register renaming, out-of-order issue/execution core architecture may implement the pipeline **2700** as follows: 1) the instruction fetch **2738** performs the fetch and length decoding stages **2702** and **2704**; 2) the decode unit **2740** performs the decode stage **2706**; 3) the rename/allocator unit **2752** performs the allocation stage **2708** and renaming stage **2710**; 4) the scheduler unit(s) **2756** performs the schedule stage **2712**; 5) the physical register file(s) unit(s) **2758** and the memory unit **2770** perform the register read/memory read stage **2714**; the execution cluster **2760** perform the execute stage **2716**; 6) the memory unit **2770** and the physical register file(s) unit(s) **2758** perform the write back/memory write stage **2718**; 7) various units may be involved in the exception handling stage **2722**; and 8) the retirement unit **2754** and the physical register file(s) unit(s) **2758** perform the commit stage **2724**.

(286) The core **2790** may support one or more instructions sets (e.g., the x86 instruction set (with some extensions that have been added with newer versions); the MIPS instruction set of MIPS Technologies of Sunnyvale, CA; the ARM instruction set (with optional additional extensions such as NEON) of ARM Holdings of Sunnyvale, CA), including the instruction(s) described herein. In one embodiment, the core **2790** includes logic to support a packed data instruction set extension (e.g., AVX1, AVX2), thereby allowing the operations used by many multimedia applications to be performed using packed data.

(287) It should be understood that the core may support multithreading (executing two or more parallel sets of operations or threads), and may do so in a variety of ways including time sliced multithreading, simultaneous multithreading (where a single physical core provides a logical core for each of the threads that physical core is simultaneously multithreading), or a combination

thereof (e.g., time sliced fetching and decoding and simultaneous multithreading thereafter such as in the Intel® Hyper-Threading technology).

(288) While register renaming is described in the context of out-of-order execution, it should be understood that register renaming may be used in an in-order architecture. While the illustrated embodiment of the processor also includes separate instruction and data cache units **2734/2774** and a shared L2 cache unit **2776**, alternative embodiments may have a single internal cache for both instructions and data, such as, for example, a Level 1 (L1) internal cache, or multiple levels of internal cache. In some embodiments, the system may include a combination of an internal cache and an external cache that is external to the core and/or the processor. Alternatively, all of the cache may be external to the core and/or the processor.

(289) Specific Exemplary In-Order Core Architecture

(290) FIGS. **28A-B** illustrate a block diagram of a more specific exemplary in-order core architecture, which core would be one of several logic blocks (including other cores of the same type and/or different types) in a chip. The logic blocks communicate through a high-bandwidth interconnect network (e.g., a ring network) with some fixed function logic, memory I/O interfaces, and other necessary I/O logic, depending on the application.

(291) FIG. **28A** is a block diagram of a single processor core, along with its connection to the on-die interconnect network **2802** and with its local subset of the Level 2 (L2) cache **2804**, according to embodiments of the disclosure. In one embodiment, an instruction decode unit **2800** supports the x86 instruction set with a packed data instruction set extension. An L1 cache **2806** allows low-latency accesses to cache memory into the scalar and vector units. While in one embodiment (to simplify the design), a scalar unit **2808** and a vector unit **2810** use separate register sets (respectively, scalar registers **2812** and vector registers **2814**) and data transferred between them is written to memory and then read back in from a level 1 (L1) cache **2806**, alternative embodiments of the disclosure may use a different approach (e.g., use a single register set or include a communication path that allow data to be transferred between the two register files without being written and read back).

(292) The local subset of the L2 cache **2804** is part of a global L2 cache that is divided into separate local subsets, one per processor core. Each processor core has a direct access path to its own local subset of the L2 cache **2804**. Data read by a processor core is stored in its L2 cache subset **2804** and can be accessed quickly, in parallel with other processor cores accessing their own local L2 cache subsets. Data written by a processor core is stored in its own L2 cache subset **2804** and is flushed from other subsets, if necessary. The ring network ensures coherency for shared data. The ring network is bi-directional to allow agents such as processor cores, L2 caches and other logic blocks to communicate with each other within the chip. Each ring data-path is 1012-bits wide per direction.

(293) FIG. **28B** is an expanded view of part of the processor core in FIG. **28A** according to embodiments of the disclosure. FIG. **28B** includes an L1 data cache **2806A** part of the L1 cache **2804**, as well as more detail regarding the vector unit **2810** and the vector registers **2814**. Specifically, the vector unit **2810** is a 16-wide vector processing unit (VPU) (see the 16-wide ALU **2828**), which executes one or more of integer, single-precision float, and double-precision float instructions. The VPU supports swizzling the register inputs with swizzle unit **2820**, numeric conversion with numeric convert units **2822A-B**, and replication with replication unit **2824** on the memory input. Write mask registers **2826** allow predicating resulting vector writes.

(294) FIG. **29** is a block diagram of a processor **2900** that may have more than one core, may have an integrated memory controller, and may have integrated graphics according to embodiments of the disclosure. The solid lined boxes in FIG. **29** illustrate a processor **2900** with a single core **2902A**, a system agent **2910**, a set of one or more bus controller units **2916**, while the optional addition of the dashed lined boxes illustrates an alternative processor **2900** with multiple cores **2902A-N**, a set of one or more integrated memory controller unit(s) **2914** in the system agent unit

2910, and special purpose logic **2908**.

(295) Thus, different implementations of the processor **2900** may include: 1) a CPU with the special purpose logic **2908** being integrated graphics and/or scientific (throughput) logic (which may include one or more cores), and the cores **2902A-N** being one or more general purpose cores (e.g., general purpose in-order cores, general purpose out-of-order cores, a combination of the two); 2) a coprocessor with the cores **2902A-N** being a large number of special purpose cores intended primarily for graphics and/or scientific (throughput); and 3) a coprocessor with the cores **2902A-N** being a large number of general purpose in-order cores. Thus, the processor **2900** may be a general-purpose processor, coprocessor or special-purpose processor, such as, for example, a network or communication processor, compression engine, graphics processor, GPGPU (general purpose graphics processing unit), a high-throughput many integrated core (MIC) coprocessor (including 30 or more cores), embedded processor, or the like. The processor may be implemented on one or more chips. The processor **2900** may be a part of and/or may be implemented on one or more substrates using any of a number of process technologies, such as, for example, BiCMOS, CMOS, or NMOS.

(296) The memory hierarchy includes one or more levels of cache within the cores, a set or one or more shared cache units **2906**, and external memory (not shown) coupled to the set of integrated memory controller units **2914**. The set of shared cache units **2906** may include one or more mid-level caches, such as level 2 (L2), level 3 (L3), level 4 (L4), or other levels of cache, a last level cache (LLC), and/or combinations thereof. While in one embodiment a ring based interconnect unit **2912** interconnects the integrated graphics logic **2908**, the set of shared cache units **2906**, and the system agent unit **2910**/integrated memory controller unit(s) **2914**, alternative embodiments may use any number of well-known techniques for interconnecting such units. In one embodiment, coherency is maintained between one or more cache units **2906** and cores **2902-A-N**.

(297) In some embodiments, one or more of the cores **2902A-N** are capable of multithreading. The system agent **2910** includes those components coordinating and operating cores **2902A-N**. The system agent unit **2910** may include for example a power control unit (PCU) and a display unit. The PCU may be or include logic and components needed for regulating the power state of the cores **2902A-N** and the integrated graphics logic **2908**. The display unit is for driving one or more externally connected displays.

(298) The cores **2902A-N** may be homogenous or heterogeneous in terms of architecture instruction set; that is, two or more of the cores **2902A-N** may be capable of execution the same instruction set, while others may be capable of executing only a subset of that instruction set or a different instruction set.

(299) Exemplary Computer Architectures

(300) FIGS. **30-33** are block diagrams of exemplary computer architectures. Other system designs and configurations known in the arts for laptops, desktops, handheld PCs, personal digital assistants, engineering workstations, servers, network devices, network hubs, switches, embedded processors, digital signal processors (DSPs), graphics devices, video game devices, set-top boxes, micro controllers, cell phones, portable media players, hand held devices, and various other electronic devices, are also suitable. In general, a huge variety of systems or electronic devices capable of incorporating a processor and/or other execution logic as disclosed herein are generally suitable.

(301) Referring now to FIG. **30**, shown is a block diagram of a system **3000** in accordance with one embodiment of the present disclosure. The system **3000** may include one or more processors **3010**, **3015**, which are coupled to a controller hub **3020**. In one embodiment the controller hub **3020** includes a graphics memory controller hub (GMCH) **3090** and an Input/Output Hub (IOH) **3050** (which may be on separate chips); the GMCH **3090** includes memory and graphics controllers to which are coupled memory **3040** and a coprocessor **3045**; the IOH **3050** is couples input/output (I/O) devices **3060** to the GMCH **3090**. Alternatively, one or both of the memory and graphics

controllers are integrated within the processor (as described herein), the memory **3040** and the coprocessor **3045** are coupled directly to the processor **3010**, and the controller hub **3020** in a single chip with the IOH **3050**. Memory **3040** may include a hashing module **3040A**, for example, to store code that when executed causes a processor to perform any method of this disclosure. (302) The optional nature of additional processors **3015** is denoted in FIG. **30** with broken lines. Each processor **3010**, **3015** may include one or more of the processing cores described herein and may be some version of the processor **2900**.

(303) The memory **3040** may be, for example, dynamic random access memory (DRAM), phase change memory (PCM), or a combination of the two. For at least one embodiment, the controller hub **3020** communicates with the processor(s) **3010**, **3015** via a multi-drop bus, such as a frontside bus (FSB), point-to-point interface such as Quickpath Interconnect (QPI), or similar connection **3095**.

(304) In one embodiment, the coprocessor **3045** is a special-purpose processor, such as, for example, a high-throughput MIC processor, a network or communication processor, compression engine, graphics processor, GPGPU, embedded processor, or the like. In one embodiment, controller hub **3020** may include an integrated graphics accelerator.

(305) There can be a variety of differences between the physical resources **3010**, **3015** in terms of a spectrum of metrics of merit including architectural, microarchitectural, thermal, power consumption characteristics, and the like.

(306) In one embodiment, the processor **3010** executes instructions that control data processing operations of a general type. Embedded within the instructions may be coprocessor instructions. The processor **3010** recognizes these coprocessor instructions as being of a type that should be executed by the attached coprocessor **3045**. Accordingly, the processor **3010** issues these coprocessor instructions (or control signals representing coprocessor instructions) on a coprocessor bus or other interconnect, to coprocessor **3045**. Coprocessor(s) **3045** accept and execute the received coprocessor instructions.

(307) Referring now to FIG. **31**, shown is a block diagram of a first more specific exemplary system **3100** in accordance with an embodiment of the present disclosure. As shown in FIG. **31**, multiprocessor system **3100** is a point-to-point interconnect system, and includes a first processor **3170** and a second processor **3180** coupled via a point-to-point interconnect **3150**. Each of processors **3170** and **3180** may be some version of the processor **2900**. In one embodiment of the disclosure, processors **3170** and **3180** are respectively processors **3010** and **3015**, while coprocessor **3138** is coprocessor **3045**. In another embodiment, processors **3170** and **3180** are respectively processor **3010** coprocessor **3045**.

(308) Processors **3170** and **3180** are shown including integrated memory controller (IMC) units **3172** and **3182**, respectively. Processor **3170** also includes as part of its bus controller units point-to-point (P-P) interfaces **3176** and **3178**; similarly, second processor **3180** includes P-P interfaces **3186** and **3188**. Processors **3170**, **3180** may exchange information via a point-to-point (P-P) interface **3150** using P-P interface circuits **3178**, **3188**. As shown in FIG. **31**, IMCs **3172** and **3182** couple the processors to respective memories, namely a memory **3132** and a memory **3134**, which may be portions of main memory locally attached to the respective processors.

(309) Processors **3170**, **3180** may each exchange information with a chipset **3190** via individual P-P interfaces **3152**, **3154** using point to point interface circuits **3176**, **3194**, **3186**, **3198**. Chipset **3190** may optionally exchange information with the coprocessor **3138** via a high-performance interface **3139**. In one embodiment, the coprocessor **3138** is a special-purpose processor, such as, for example, a high-throughput MIC processor, a network or communication processor, compression engine, graphics processor, GPGPU, embedded processor, or the like.

(310) A shared cache (not shown) may be included in either processor or outside of both processors, yet connected with the processors via P-P interconnect, such that either or both processors' local cache information may be stored in the shared cache if a processor is placed into a

low power mode.

(311) Chipset **3190** may be coupled to a first bus **3116** via an interface **3196**. In one embodiment, first bus **3116** may be a Peripheral Component Interconnect (PCI) bus, or a bus such as a PCI Express bus or another third generation I/O interconnect bus, although the scope of the present disclosure is not so limited.

(312) As shown in FIG. **31**, various I/O devices **3114** may be coupled to first bus **3116**, along with a bus bridge **3118** which couples first bus **3116** to a second bus **3120**. In one embodiment, one or more additional processor(s) **3115**, such as coprocessors, high-throughput MIC processors, GPGPU's, accelerators (such as, e.g., graphics accelerators or digital signal processing (DSP) units), field programmable gate arrays, or any other processor, are coupled to first bus **3116**. In one embodiment, second bus **3120** may be a low pin count (LPC) bus. Various devices may be coupled to a second bus **3120** including, for example, a keyboard and/or mouse **3122**, communication devices **3127** and a storage unit **3128** such as a disk drive or other mass storage device which may include instructions/code and data **3130**, in one embodiment. Further, an audio I/O **3124** may be coupled to the second bus **3120**. Note that other architectures are possible. For example, instead of the point-to-point architecture of FIG. **31**, a system may implement a multi-drop bus or other such architecture.

(313) Referring now to FIG. **32**, shown is a block diagram of a second more specific exemplary system **3200** in accordance with an embodiment of the present disclosure. Like elements in FIGS. **31** and **32** bear like reference numerals, and certain aspects of FIG. **31** have been omitted from FIG. **32** in order to avoid obscuring other aspects of FIG. **32**.

(314) FIG. **32** illustrates that the processors **3170**, **3180** may include integrated memory and I/O control logic ("CL") **3172** and **3182**, respectively. Thus, the CL **3172**, **3182** include integrated memory controller units and include I/O control logic. FIG. **32** illustrates that not only are the memories **3132**, **3134** coupled to the CL **3172**, **3182**, but also that I/O devices **3214** are also coupled to the control logic **3172**, **3182**. Legacy I/O devices **3215** are coupled to the chipset **3190**.

(315) Referring now to FIG. **33**, shown is a block diagram of a SoC **3300** in accordance with an embodiment of the present disclosure. Similar elements in FIG. **29** bear like reference numerals. Also, dashed lined boxes are optional features on more advanced SoCs. In FIG. **33**, an interconnect unit(s) **3302** is coupled to: an application processor **3310** which includes a set of one or more cores **202A-N** and shared cache unit(s) **2906**; a system agent unit **2910**; a bus controller unit(s) **2916**; an integrated memory controller unit(s) **2914**; a set of one or more coprocessors **3320** which may include integrated graphics logic, an image processor, an audio processor, and a video processor; a static random access memory (SRAM) unit **3330**; a direct memory access (DMA) unit **3332**; and a display unit **3340** for coupling to one or more external displays. In one embodiment, the coprocessor(s) **3320** include a special-purpose processor, such as, for example, a network or communication processor, compression engine, GPGPU, a high-throughput MIC processor, embedded processor, or the like.

(316) Embodiments (e.g., of the mechanisms) disclosed herein may be implemented in hardware, software, firmware, or a combination of such implementation approaches. Embodiments of the disclosure may be implemented as computer programs or program code executing on programmable systems comprising at least one processor, a storage system (including volatile and non-volatile memory and/or storage elements), at least one input device, and at least one output device.

(317) Program code, such as code **3130** illustrated in FIG. **31**, may be applied to input instructions to perform the functions described herein and generate output information. The output information may be applied to one or more output devices, in known fashion. For purposes of this application, a processing system includes any system that has a processor, such as, for example; a digital signal processor (DSP), a microcontroller, an application specific integrated circuit (ASIC), or a microprocessor.

(318) The program code may be implemented in a high level procedural or object oriented programming language to communicate with a processing system. The program code may also be implemented in assembly or machine language, if desired. In fact, the mechanisms described herein are not limited in scope to any particular programming language. In any case, the language may be a compiled or interpreted language.

(319) One or more aspects of at least one embodiment may be implemented by representative instructions stored on a machine-readable medium which represents various logic within the processor, which when read by a machine causes the machine to fabricate logic to perform the techniques described herein. Such representations, known as “IP cores,” may be stored on a tangible, machine readable medium and supplied to various customers or manufacturing facilities to load into the fabrication machines that actually make the logic or processor.

(320) Such machine-readable storage media may include, without limitation, non-transitory, tangible arrangements of articles manufactured or formed by a machine or device, including storage media such as hard disks, any other type of disk including floppy disks, optical disks, compact disk read-only memories (CD-ROMs), compact disk rewritables (CD-RWs), and magneto-optical disks, semiconductor devices such as read-only memories (ROMs), random access memories (RAMs) such as dynamic random access memories (DRAMs), static random access memories (SRAMs), erasable programmable read-only memories (EPROMs), flash memories, electrically erasable programmable read-only memories (EEPROMs), phase change memory (PCM), magnetic or optical cards, or any other type of media suitable for storing electronic instructions.

(321) Accordingly, embodiments of the disclosure also include non-transitory, tangible machine-readable media containing instructions or containing design data, such as Hardware Description Language (HDL), which defines structures, circuits, apparatuses, processors and/or system features described herein. Such embodiments may also be referred to as program products.

(322) Emulation (Including Binary Translation, Code Morphing, Etc.)

(323) In some cases, an instruction converter may be used to convert an instruction from a source instruction set to a target instruction set. For example, the instruction converter may translate (e.g., using static binary translation, dynamic binary translation including dynamic compilation), morph, emulate, or otherwise convert an instruction to one or more other instructions to be processed by the core. The instruction converter may be implemented in software, hardware, firmware, or a combination thereof. The instruction converter may be on processor, off processor, or part on and part off processor.

(324) FIG. 34 is a block diagram contrasting the use of a software instruction converter to convert binary instructions in a source instruction set to binary instructions in a target instruction set according to embodiments of the disclosure. In the illustrated embodiment, the instruction converter is a software instruction converter, although alternatively the instruction converter may be implemented in software, firmware, hardware, or various combinations thereof. FIG. 34 shows a program in a high level language 3402 may be compiled using an x86 compiler 3404 to generate x86 binary code 3406 that may be natively executed by a processor with at least one x86 instruction set core 3416. The processor with at least one x86 instruction set core 3416 represents any processor that can perform substantially the same functions as an Intel® processor with at least one x86 instruction set core by compatibly executing or otherwise processing (1) a substantial portion of the instruction set of the Intel® x86 instruction set core or (2) object code versions of applications or other software targeted to run on an Intel® processor with at least one x86 instruction set core, in order to achieve substantially the same result as an Intel® processor with at least one x86 instruction set core. The x86 compiler 3404 represents a compiler that is operable to generate x86 binary code 3406 (e.g., object code) that can, with or without additional linkage processing, be executed on the processor with at least one x86 instruction set core 3416. Similarly, FIG. 34 shows the program in the high level language 3402 may be compiled using an alternative

instruction set compiler **3408** to generate alternative instruction set binary code **3410** that may be natively executed by a processor without at least one x86 instruction set core **3414** (e.g., a processor with cores that execute the MIPS instruction set of MIPS Technologies of Sunnyvale, CA and/or that execute the ARM instruction set of ARM Holdings of Sunnyvale, CA). The instruction converter **3412** is used to convert the x86 binary code **3406** into code that may be natively executed by the processor without an x86 instruction set core **3414**. This converted code is not likely to be the same as the alternative instruction set binary code **3410** because an instruction converter capable of this is difficult to make; however, the converted code will accomplish the general operation and be made up of instructions from the alternative instruction set. Thus, the instruction converter **3412** represents software, firmware, hardware, or a combination thereof that, through emulation, simulation or any other process, allows a processor or other electronic device that does not have an x86 instruction set processor or core to execute the x86 binary code **3406**.

Claims

1. A processor comprising: a fetch circuit to fetch a single instruction, the single instruction having at least one first field to identify eight 32-bit state elements A, B, C, D, E, F, G, and H for a round according to an SM3 hash function and at least one second field to identify an input message; a decode circuit to decode the single instruction; and an execution circuit to perform operations corresponding to the single instruction, including to: rotate the state element C left by 9 bits to form a rotated state element C, rotate the state element D left by 9 bits to form a rotated state element D, rotate the state element G left by 19 bits to form a rotated state element G, rotate the state element H left by 19 bits to form a rotated state element H, perform two rounds according to the SM3 hash function on the input message and the state element A, the state element B, the rotated state element C, the rotated state element D, the state element E, the state element F, the rotated state element G, and the rotated state element H to generate an updated state element A, an updated state element B, an updated state element E, and an updated state element F, and store the updated state element A, the updated state element B, the updated state element E, and the updated state element F to a location specified by the single instruction.
2. The processor of claim 1, wherein the at least one first field includes: a first field to identify a first register that is to store the state elements C, D, G, and H, and a second field to identify a second register that is to store the state elements A, B, E, and F.
3. The processor of claim 2, wherein the at least one second field includes a third field to identify the input message.
4. The processor of claim 2, wherein the location specified by the single instruction to store the updated state element A, the updated state element B, the updated state element E, and the updated state element F is the first register.
5. The processor of claim 1, wherein the state elements A, B, E, and F are to be stored in a first 128-bit register.
6. The processor of claim 5, wherein the state element A is to be stored in a least significant doubleword data element position of the first 128-bit register.
7. The processor of claim 6, wherein the state elements C, D, G, and H are to be stored in a second 128-bit register.
8. The processor of claim 7, wherein the state element C is to be stored in a least significant doubleword data element position of the second 128-bit register.
9. A method comprising: fetching, by a fetch circuit, a single instruction, the single instruction having at least one first field to identify eight 32-bit state elements A, B, C, D, E, F, G, and H for a round according to an SM3 hash function and at least one second field to identify an input message; decoding, by a decode circuit, the single instruction; and executing, by an execution circuit, the single instruction to perform operations comprising: rotating the state element C left by

9 bits to form a rotated state element C, rotating the state element D left by 9 bits to form a rotated state element D, rotating the state element G left by 19 bits to form a rotated state element G, rotating the state element H left by 19 bits to form a rotated state element H, performing two rounds according to the SM3 hash function on the input message and the state element A, the state element B, the rotated state element C, the rotated state element D, the state element E, the state element F, the rotated state element G, and the rotated state element H to generate an updated state element A, an updated state element B, an updated state element E, and an updated state element F, and storing the updated state element A, the updated state element B, the updated state element E, and the updated state element F to a location specified by the single instruction.

10. The method of claim 9, wherein the at least one first field includes: a first field to identify a first register that is to store the state elements C, D, G, and H, and a second field to identify a second register that is to store the state elements A, B, E, and F.

11. The method of claim 10, wherein the at least one second field includes a third field to identify the input message.

12. The method of claim 10, wherein the location specified by the single instruction that stores the updated state element A, the updated state element B, the updated state element E, and the updated state element F is the first register.

13. The method of claim 9, wherein the state elements A, B, E, and F are stored in a first 128-bit register.

14. The method of claim 13, wherein the state element A is stored in a least significant doubleword data element position of the first 128-bit register.

15. The method of claim 14, wherein the state elements C, D, G, and H are stored in a second 128-bit register.

16. The method of claim 15, wherein the state element C is stored in a least significant doubleword data element position of the second 128-bit register.

17. A non-transitory machine readable medium that stores code that when executed by a machine causes the machine to perform a method comprising: fetching, by a fetch circuit, a single instruction, the single instruction having at least one first field to identify eight 32-bit state elements A, B, C, D, E, F, G, and H for a round according to an SM3 hash function and at least one second field to identify an input message; decoding, by a decode circuit, the single instruction; and executing, by an execution circuit, the single instruction to perform operations comprising: rotating the state element C left by 9 bits to form a rotated state element C, rotating the state element D left by 9 bits to form a rotated state element D, rotating the state element G left by 19 bits to form a rotated state element G, rotating the state element H left by 19 bits to form a rotated state element H, performing two rounds according to the SM3 hash function on the input message and the state element A, the state element B, the rotated state element C, the rotated state element D, the state element E, the state element F, the rotated state element G, and the rotated state element H to generate an updated state element A, an updated state element B, an updated state element E, and an updated state element F, and storing the updated state element A, the updated state element B, the updated state element E, and the updated state element F to a location specified by the single instruction.

18. The non-transitory machine readable medium of claim 17, wherein the at least one first field includes: a first field to identify a first register that is to store the state elements C, D, G, and H, and a second field to identify a second register that is to store the state elements A, B, E, and F.

19. The non-transitory machine readable medium of claim 18, wherein the at least one second field includes a third field to identify the input message.

20. The non-transitory machine readable medium of claim 18, wherein the location specified by the single instruction that stores the updated state element A, the updated state element B, the updated state element E, and the updated state element F is the first register.

21. The non-transitory machine readable medium of claim 17, wherein the state elements A, B, E,

and F are stored in a first 128-bit register.

22. The non-transitory machine readable medium of claim 21, wherein the state element A is stored in a least significant doubleword data element position of the first 128-bit register.

23. The non-transitory machine readable medium of claim 22, wherein the state elements C, D, G, and H are stored in a second 128-bit register.

24. The non-transitory machine readable medium of claim 23, wherein the state element C is stored in a least significant doubleword data element position of the second 128-bit register.
