

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12399742
Kind Code	B2
Date of Patent	August 26, 2025
Inventor(s)	Singh; Naorem Khogendro et al.

Communication and synchronization with edge systems

Abstract

A scalable Internet of Things (IoT) system may include multiple instances of an IoT manager, each instance respectively configured to connect to a respective edge system of multiple edge systems. The IoT system may further include a containerized system configured to allow any instance of the IoT manager to deploy data pipelines to any edge system of the multiple edge systems in delta communications. Any instance of the IoT manager may send a change message to any edge system via a publish/subscribe notification method. In some examples, a centralized IoT manager may form a secure communication with an edge system, synchronize an object model with an edge object model for the edge system, and maintain the edge system using delta change communications. The IoT system may facilitate any instance of the IoT manager to subscribe a communication channel with an associated edge system for receiving update notification.

Inventors:	Singh; Naorem Khogendro (San Jose, CA), Goli; Sandeep Reddy (San Jose, CA), Vaghani; Satyam B. (San Jose, CA), Perng; Shyan-Ming (San Jose, CA)
Applicant:	Nutanix, Inc. (San Jose, CA)
Family ID:	1000008779981
Assignee:	Nutanix, Inc. (San Jose, CA)
Appl. No.:	18/672892
Filed:	May 23, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20250060999 A1	Feb. 20, 2025

Related U.S. Application Data

continuation parent-doc US 18185300 20230316 US 12026551 child-doc US 18672892
division parent-doc US 16666242 20191028 US 11635990 20230425 child-doc US 18185300
us-provisional-application US 62869348 20190701

Publication Classification

Int. Cl.: G06F9/54 (20060101); G06F9/38 (20180101); G06F9/50 (20060101); H04L67/125 (20220101)

U.S. Cl.:

CPC G06F9/5011 (20130101); G06F9/3869 (20130101); G06F9/5077 (20130101); G06F9/542 (20130101); H04L67/125 (20130101);

Field of Classification Search

CPC: G06F (9/4812); G06F (9/54); G06F (9/542); G06F (9/544); G06F (9/546); H04L (29/06)

References Cited

U.S. PATENT DOCUMENTS				
Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC

6550057	12/2002	Bowman-amuah	N/A	N/A
7237360	12/2006	Moncho et al.	N/A	N/A
7310664	12/2006	Merchant et al.	N/A	N/A
8543665	12/2012	Ansari et al.	N/A	N/A
8549518	12/2012	Aron et al.	N/A	N/A
8601473	12/2012	Aron et al.	N/A	N/A
8850130	12/2013	Aron et al.	N/A	N/A
8863124	12/2013	Aron	N/A	N/A
9009106	12/2014	Aron et al.	N/A	N/A
9069708	12/2014	Gill et al.	N/A	N/A
9210534	12/2014	Matthieu et al.	N/A	N/A
9244951	12/2015	Mandelstein et al.	N/A	N/A
9253252	12/2015	Agarwal et al.	N/A	N/A
9336132	12/2015	Aron et al.	N/A	N/A
9600494	12/2016	Maluf et al.	N/A	N/A
9606794	12/2016	Chou et al.	N/A	N/A
9633197	12/2016	Lakshmanan et al.	N/A	N/A
9634893	12/2016	Boutros et al.	N/A	N/A
9641650	12/2016	Virkki et al.	N/A	N/A
9652265	12/2016	Narayanasamy et al.	N/A	N/A
9684502	12/2016	Fu et al.	N/A	N/A
9729411	12/2016	Purusothaman	N/A	N/A
9736194	12/2016	Rao et al.	N/A	N/A
9747287	12/2016	Bhardwaj et al.	N/A	N/A
9769032	12/2016	Ghodsí et al.	N/A	N/A
9772866	12/2016	Aron et al.	N/A	N/A
9836296	12/2016	Vandikas et al.	N/A	N/A
9860677	12/2017	Agerstam et al.	N/A	N/A
9917865	12/2017	Arora et al.	N/A	N/A
9977415	12/2017	Zimmerman et al.	N/A	N/A
10091270	12/2017	Fang	N/A	N/A
10149154	12/2017	Zimmerman et al.	N/A	N/A
10156842	12/2017	Wu et al.	N/A	N/A
10181978	12/2018	Argenti	N/A	N/A
10225335	12/2018	Fu et al.	N/A	N/A
10230798	12/2018	Doraiswamy et al.	N/A	N/A
10262019	12/2018	Reiner et al.	N/A	N/A
10291714	12/2018	Mathews et al.	N/A	N/A
10306513	12/2018	Bartfai-Walcott et al.	N/A	N/A
10489138	12/2018	Wu et al.	N/A	N/A
10515119	12/2018	Kirk et al.	N/A	N/A
10534629	12/2019	St. Pierre et al.	N/A	N/A
10567925	12/2019	Ly et al.	N/A	N/A
10652226	12/2019	Islam et al.	N/A	N/A
10685295	12/2019	Ross et al.	N/A	N/A
10726302	12/2019	Zhang	N/A	G06V 10/764
10769113	12/2019	Jones	N/A	N/A
10838833	12/2019	Jibaja et al.	N/A	N/A
10884808	12/2020	Chintalapally	N/A	G06F 9/5044
10893116	12/2020	Koehler	N/A	N/A
10924342	12/2020	Joshi et al.	N/A	N/A
11108849	12/2020	Tofighbakhsh	N/A	H04L 41/5025
11178213	12/2020	Desai et al.	N/A	N/A
11194483	12/2020	Dontu et al.	N/A	N/A
11301762	12/2021	Chen et al.	N/A	N/A
11316733	12/2021	Johson et al.	N/A	N/A
11379525	12/2021	Deutsch et al.	N/A	N/A
11397692	12/2021	Chhabra et al.	N/A	N/A
11409756	12/2021	Park et al.	N/A	N/A
11422863	12/2021	Sengupta et al.	N/A	N/A
11429893	12/2021	Tong	N/A	N/A
11501881	12/2021	Patil et al.	N/A	N/A
11635990	12/2022	Singh et al.	N/A	N/A
11665221	12/2022	Joshi et al.	N/A	N/A
11726764	12/2022	Joshi et al.	N/A	N/A
12021915	12/2023	Joshi et al.	N/A	N/A
12026551	12/2023	Singh et al.	N/A	N/A
12155731	12/2023	Joshi et al.	N/A	N/A
12159178	12/2023	Patil et al.	N/A	N/A

2002/0143819	12/2001	Han et al.	N/A	N/A
2004/0177359	12/2003	Bauch et al.	N/A	N/A
2005/0060328	12/2004	Suhonen et al.	N/A	N/A
2008/0270515	12/2007	Chen et al.	N/A	N/A
2009/0204711	12/2008	Binyamin	N/A	N/A
2009/0260004	12/2008	Datta et al.	N/A	N/A
2010/0008510	12/2009	Zayas	N/A	N/A
2010/0077473	12/2009	Takeshita et al.	N/A	N/A
2010/0175061	12/2009	Maeda et al.	N/A	N/A
2011/0153824	12/2010	Chikando et al.	N/A	N/A
2011/0231899	12/2010	Pulier et al.	N/A	N/A
2011/0251992	12/2010	Bethlehem et al.	N/A	N/A
2011/0265077	12/2010	Collison et al.	N/A	N/A
2012/0102486	12/2011	Yendluri	N/A	N/A
2012/0167081	12/2011	Sedayao et al.	N/A	N/A
2012/0265884	12/2011	Zhang et al.	N/A	N/A
2012/0266156	12/2011	Spivak et al.	N/A	N/A
2012/0281708	12/2011	Chauhan et al.	N/A	N/A
2013/0219479	12/2012	Desoto et al.	N/A	N/A
2013/0227573	12/2012	Morsi et al.	N/A	N/A
2013/0304903	12/2012	Mick et al.	N/A	N/A
2013/0332916	12/2012	Chinn et al.	N/A	N/A
2014/0074539	12/2013	Doering et al.	N/A	N/A
2014/0075412	12/2013	Kannan et al.	N/A	N/A
2014/0075431	12/2013	Kumar et al.	N/A	N/A
2014/0164486	12/2013	Ravinchandran et al.	N/A	N/A
2014/0279899	12/2013	Gu et al.	N/A	N/A
2014/0317261	12/2013	Shatzkamer et al.	N/A	N/A
2014/0330948	12/2013	Dunn et al.	N/A	N/A
2014/0372508	12/2013	Fausak et al.	N/A	N/A
2015/0057817	12/2014	Endrizzi et al.	N/A	N/A
2015/0067030	12/2014	Smith et al.	N/A	N/A
2015/0074106	12/2014	Ji	N/A	N/A
2015/0120893	12/2014	Sapaliga et al.	N/A	N/A
2015/0261876	12/2014	Trikha et al.	N/A	N/A
2015/0347542	12/2014	Sullivan et al.	N/A	N/A
2016/0007138	12/2015	Palanisamy et al.	N/A	N/A
2016/0092176	12/2015	Straub et al.	N/A	N/A
2016/0092180	12/2015	Straub	N/A	N/A
2016/0092348	12/2015	Straub et al.	N/A	N/A
2016/0094398	12/2015	Choudhary et al.	N/A	N/A
2016/0098265	12/2015	Mahajan et al.	N/A	N/A
2016/0112268	12/2015	Chung et al.	N/A	N/A
2016/0156614	12/2015	Jain et al.	N/A	N/A
2016/0197830	12/2015	Ulevitch et al.	N/A	N/A
2016/0202964	12/2015	Butcher et al.	N/A	N/A
2016/0216959	12/2015	Kurian et al.	N/A	N/A
2016/0234071	12/2015	Nambiar et al.	N/A	N/A
2016/0315848	12/2015	Weinstein	N/A	N/A
2016/0315912	12/2015	Mayya et al.	N/A	N/A
2016/0323161	12/2015	Cuervo Laffaye et al.	N/A	N/A
2016/0323361	12/2015	Austel et al.	N/A	N/A
2016/0337104	12/2015	Kalligudd	N/A	N/A
2016/0337175	12/2015	Rao	N/A	N/A
2016/0342906	12/2015	Shaashua et al.	N/A	N/A
2016/0344745	12/2015	Johnson et al.	N/A	N/A
2016/0345516	12/2015	Britt et al.	N/A	N/A
2016/0357525	12/2015	Wee et al.	N/A	N/A
2017/0005820	12/2016	Zimmerman et al.	N/A	N/A
2017/0006135	12/2016	Siebel et al.	N/A	N/A
2017/0048079	12/2016	Nethi et al.	N/A	N/A
2017/0060574	12/2016	Malladi et al.	N/A	N/A
2017/0097950	12/2016	Meacham et al.	N/A	N/A
2017/0099176	12/2016	Jain	N/A	N/A
2017/0102931	12/2016	Risbood et al.	N/A	N/A
2017/0123939	12/2016	Maheshwari et al.	N/A	N/A
2017/0126809	12/2016	Chen et al.	N/A	N/A
2017/0142068	12/2016	Devarajan et al.	N/A	N/A
2017/0149931	12/2016	Lochhead et al.	N/A	N/A

2017/0168813	12/2016	Pogrebinsky et al.	N/A	N/A
2017/0171607	12/2016	Britt	N/A	N/A
2017/0177334	12/2016	Chou et al.	N/A	N/A
2017/0177877	12/2016	Suarez et al.	N/A	N/A
2017/0180272	12/2016	Bernath	N/A	N/A
2017/0180346	12/2016	Suarez et al.	N/A	N/A
2017/0185507	12/2016	Eberlein	N/A	N/A
2017/0185922	12/2016	Lange et al.	N/A	N/A
2017/0187807	12/2016	Clernon	N/A	N/A
2017/0201411	12/2016	Mladin et al.	N/A	N/A
2017/0201585	12/2016	Doraiswamy et al.	N/A	N/A
2017/0212718	12/2016	Nelson et al.	N/A	N/A
2017/0213157	12/2016	Bugay et al.	N/A	N/A
2017/0244600	12/2016	Hussein et al.	N/A	N/A
2017/0257432	12/2016	Fu et al.	N/A	N/A
2017/0289173	12/2016	Resch et al.	N/A	N/A
2017/0315820	12/2016	Entezari et al.	N/A	N/A
2017/0347264	12/2016	Holland et al.	N/A	N/A
2018/0007055	12/2017	Infante-Lopez et al.	N/A	N/A
2018/0013819	12/2017	Li	N/A	N/A
2018/0034914	12/2017	Christopher et al.	N/A	N/A
2018/0054315	12/2017	Liu et al.	N/A	N/A
2018/0054490	12/2017	Wadhwa et al.	N/A	N/A
2018/0067830	12/2017	Jagtiani et al.	N/A	N/A
2018/0092151	12/2017	Liu et al.	N/A	N/A
2018/0101415	12/2017	Mahindru et al.	N/A	N/A
2018/0109395	12/2017	Berdy et al.	N/A	N/A
2018/0109650	12/2017	Berdy et al.	N/A	N/A
2018/0109929	12/2017	Ly et al.	N/A	N/A
2018/0123820	12/2017	Kim	N/A	N/A
2018/0159745	12/2017	Byers et al.	N/A	N/A
2018/0167392	12/2017	Zakaria	N/A	N/A
2018/0212970	12/2017	Chen et al.	N/A	N/A
2018/0219877	12/2017	Hsu et al.	N/A	N/A
2018/0234351	12/2017	Amento et al.	N/A	N/A
2018/0293463	12/2017	Brown	N/A	N/A
2018/0295194	12/2017	Deraz et al.	N/A	N/A
2018/0300124	12/2017	Malladi et al.	N/A	N/A
2018/0302266	12/2017	Makovsky et al.	N/A	N/A
2018/0307464	12/2017	Bijani et al.	N/A	N/A
2018/0307537	12/2017	Chen et al.	N/A	N/A
2018/0309819	12/2017	Thompson	N/A	N/A
2018/0314808	12/2017	Casey et al.	N/A	N/A
2018/0324204	12/2017	McClory et al.	N/A	N/A
2018/0332116	12/2017	George et al.	N/A	N/A
2018/0338242	12/2017	Li et al.	N/A	N/A
2018/0373419	12/2017	Chen et al.	N/A	N/A
2018/0373555	12/2017	Gupta et al.	N/A	N/A
2019/0014048	12/2018	Krishna Singuru	N/A	N/A
2019/0018951	12/2018	James et al.	N/A	N/A
2019/0026082	12/2018	Shalev et al.	N/A	N/A
2019/0034837	12/2018	Lou et al.	N/A	N/A
2019/0041824	12/2018	Chavez et al.	N/A	N/A
2019/0045033	12/2018	Agerstam et al.	N/A	N/A
2019/0068445	12/2018	Chauhan et al.	N/A	N/A
2019/0087220	12/2018	Turner	N/A	N/A
2019/0098113	12/2018	Park et al.	N/A	N/A
2019/0102265	12/2018	Ngo et al.	N/A	N/A
2019/0109816	12/2018	Liu et al.	N/A	N/A
2019/0114211	12/2018	Reddipalli et al.	N/A	N/A
2019/0121889	12/2018	Gold et al.	N/A	N/A
2019/0123959	12/2018	Joshi et al.	N/A	N/A
2019/0141022	12/2018	Reeve et al.	N/A	N/A
2019/0146773	12/2018	Attard	N/A	N/A
2019/0158353	12/2018	Johnson et al.	N/A	N/A
2019/0158600	12/2018	Cook	N/A	N/A
2019/0182333	12/2018	Bartfai-Walcott et al.	N/A	N/A
2019/0188742	12/2018	Vasudevan et al.	N/A	N/A
2019/0190776	12/2018	Bregman et al.	N/A	N/A

2019/0213273	12/2018	Vasudevan et al.	N/A	N/A
2019/0243836	12/2018	Nanda et al.	N/A	N/A
2019/0251166	12/2018	Penrose et al.	N/A	N/A
2019/0286353	12/2018	Soni et al.	N/A	N/A
2019/0295012	12/2018	Marinescu et al.	N/A	N/A
2019/0319919	12/2018	Knecht et al.	N/A	N/A
2019/0320038	12/2018	Walsh et al.	N/A	N/A
2019/0342182	12/2018	Dhanabalan et al.	N/A	N/A
2019/0384641	12/2018	Zhao et al.	N/A	N/A
2020/0014607	12/2019	Gangaadhar et al.	N/A	N/A
2020/0014633	12/2019	You et al.	N/A	N/A
2020/0034776	12/2019	Peran et al.	N/A	N/A
2020/0073739	12/2019	Rungta et al.	N/A	N/A
2020/0092789	12/2019	Lee et al.	N/A	N/A
2020/0097274	12/2019	Sarkar et al.	N/A	N/A
2020/0104723	12/2019	Reissner et al.	N/A	N/A
2020/0112487	12/2019	Inamdar et al.	N/A	N/A
2020/0122038	12/2019	Ebrahimi et al.	N/A	N/A
2020/0126129	12/2019	Lkhamsuren	N/A	N/A
2020/0127832	12/2019	Ebrahimi	N/A	N/A
2020/0177503	12/2019	Hooda et al.	N/A	N/A
2020/0177630	12/2019	Penner et al.	N/A	N/A
2020/0192690	12/2019	Gupta et al.	N/A	N/A
2020/0218580	12/2019	Kim	N/A	N/A
2020/0250002	12/2019	Gururaj et al.	N/A	N/A
2020/0258627	12/2019	Setegn et al.	N/A	N/A
2020/0274776	12/2019	Nishikawa	N/A	N/A
2020/0287737	12/2019	Mishra et al.	N/A	N/A
2020/0336467	12/2019	Subbarayan et al.	N/A	N/A
2020/0356415	12/2019	Goli	N/A	N/A
2021/0004270	12/2020	Singh et al.	N/A	N/A
2021/0005330	12/2020	Patil et al.	N/A	N/A
2021/0006636	12/2020	Koehler et al.	N/A	N/A
2021/0042104	12/2020	Tashkandi	N/A	N/A
2021/0042160	12/2020	Alamouti et al.	N/A	N/A
2021/0044579	12/2020	Nelson-Gal et al.	N/A	N/A
2021/0058338	12/2020	Liu	N/A	N/A
2021/0067607	12/2020	Gardner et al.	N/A	N/A
2021/0084670	12/2020	Chauhan et al.	N/A	N/A
2021/0089408	12/2020	Park et al.	N/A	N/A
2021/0112059	12/2020	Heldman et al.	N/A	N/A
2021/0112128	12/2020	Joshi et al.	N/A	N/A
2021/0140815	12/2020	Pretorius et al.	N/A	N/A
2021/0160338	12/2020	Koehler et al.	N/A	N/A
2021/0232344	12/2020	Corrie	N/A	N/A
2021/0255846	12/2020	Mamgain et al.	N/A	N/A
2021/0271565	12/2020	Bhavanarushi et al.	N/A	N/A
2021/0311764	12/2020	Rosoff et al.	N/A	N/A
2021/0342193	12/2020	Anand	N/A	N/A
2021/0373965	12/2020	Hadas et al.	N/A	N/A
2021/0400043	12/2020	Su et al.	N/A	N/A
2022/0083018	12/2021	Majeed et al.	N/A	N/A
2022/0083389	12/2021	Poothia et al.	N/A	N/A
2022/0121543	12/2021	Poothia et al.	N/A	N/A
2022/0138070	12/2021	Mokashi et al.	N/A	N/A
2022/0147336	12/2021	Joshi et al.	N/A	N/A
2022/0159093	12/2021	Joshi et al.	N/A	N/A
2022/0279046	12/2021	Perng et al.	N/A	N/A
2022/0334882	12/2021	Wang	N/A	N/A
2023/0070811	12/2022	Patil et al.	N/A	N/A
2023/0141808	12/2022	Joshi et al.	N/A	N/A
2023/0289225	12/2022	Singh et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
101916200	12/2014	CN	N/A
110795442	12/2021	CN	N/A
2831746	12/2014	EP	N/A
9945465	12/1998	WO	N/A

2014007811	12/2013	WO	N/A
2020096639	12/2019	WO	N/A

OTHER PUBLICATIONS

“Architecting For The Cloud: AWS Best Practice”, Amazon Web Services https://d1.awsstatic.com/whitepapers/AWS_Cloud_Best_Practices.pdf, Oct. 2018, pp. 1-50, No Author. cited by applicant

“Architecting VMware Unified Access Gateway”, <https://www.youtube.com/watch?v=URSdJ9qCQKo&t=234s>, Apr. 17, 2019, pp. 1, No author. cited by applicant

“AWS Architecture Monthly—Internet of Things Jun. 2019”, AWS | <https://d1.awsstatic.com/whitepapers/architecture-monthly/AWS-Architecture-Monthly-Jun.2019.pdf>, Jun. 2019, No author. cited by applicant

“AWS IoT Greengrass”, <https://web.archive.org/web/20190624094650/https://aws.amazon.com/greengrass/>, Jun. 2019, pp. 1-13, No author. cited by applicant

“AWS IoT vs. Google IoT vs. Azure IoT”, Bizety | <https://www.bizety.com/2018/08/28/aws-iot-vs-google-iot-vs-azure-iot/>, Aug. 28, 2018, pp. 7, No author. cited by applicant

“Azure IoT Central intro walkthrough”, https://www.youtube.com/watch?v=G32stXSwtYA&ab_channel=MicrosoftDeveloper, Feb. 2018, No author. cited by applicant

“Backup & Secure”, USGS, 2018, No author. cited by applicant

“Best Practices for Cloud Management”, Service Now <https://www.servicenow.com/content/dam/servicenow-assets/public/en-us/doc-type/resource-center/solution-brief/sb-cloud-management.pdf>, Jun. 2017, pp. 1-2, No author. cited by applicant

“Cascade Mode Deployment”, <https://docs.vmware.com/en/Unified-Access-Gateway/3.10/com.vmware.uag-310-deploy-config.doc/GUID-AWT-DEPLOYMENT-CASCADE.html#GUID-AWT-DEPLOYMENT-CASCADE>, Jul. 2, 2019, pp. 2, No author. cited by applicant

“Citrix XenDesktop 7.1 on Microsoft Hyper-V Server 2012 R2 on Nutanix Virtual Computing Platform”, Citrix APAC Solutions, Jun. 25, 2014, pp. 1-94, No author. cited by applicant

“Cloud Iot Core”, <https://web.archive.org/web/20190129000453/https://cloud.google.com/iot-core/>, <https://web.archive.org/web/20190129000453/https://cloud.google.com/iot-core/>, Jan. 2019, pp. 1-10, No author. cited by applicant

“Cloud IoT Core—Devices, configuration, and state”, Google | <https://web.archive.org/web/20190330153113/https://cloud.google.com/iot/docs/concepts/devices>, Mar. 2019, No author. cited by applicant

“Cloud IoT Core Private Beta”, Google | <https://web.archive.org/web/20170518022234/https://cloud.google.com/iot-core/>, May 2017, No author. cited by applicant

“Cloud, Fog and Edge Computing—What's the Difference?”, <https://www.winsystems.com/cloud-fog-and-edge-computing-whats-the-difference>, Dec. 4, 2017, pp. 1-10, No author. cited by applicant

“Comprehensive Guide on Upgrading PKS”, PKS 1.3 <https://kb.vmware.com/sfc/servlet.shepherd/version/download/068f4000009EfWPAAO>, Apr. 2019, pp. 1-45, No author. cited by applicant

“Deploying and Configuring VMware Unified Access Gateway”, VMware | Unified Access Gateway 3.1 <https://techzone.vmware.com/configuring-edge-services-vmware-unified-access-gateway-vmware-workspace-one-operational-tutorial#overview>, Oct. 27, 2017, pp. 1-89, No author. cited by applicant

“Extended offline operation with Azure IoTEdge”, Microsoft | <https://azure.microsoft.com/en-us/blog/extended-offline-operation-with-azure-iot-edge/>, Sep. 2018, No author. cited by applicant

“Extract, Transform, Load with AWS IoT Greengrass Solution Accelerator”, AWS | <https://aws.amazon.com/iot/solutions/etl-accelerator/>, Oct. 2019, pp. 1-6, No author. cited by applicant

“Google Cloud IoT Core”, Google | <https://www.slideshare.net/idof/google-cloud-iot-core>, Sep. 2018, No author. cited by applicant

“Google IoT Platform: Awesome Tools for Any Project”, <https://indeema.com/blog/google-cloud>, Feb. 20, 2019, pp. 1-9, No author. cited by applicant

“How to Troubleshoot PKS Upgrade Issues”, VMware Tanzu Supprt Hub https://community.pivotal.io/s/article/how-to-troubleshoot-pks-upgrade-issues?language=en_US, Jul. 17, 2019, pp. 1-11, No author. cited by applicant

“Introducing Google Cloud IoT Core: forsecurely connecting and managing IoTdevices at scale”, Google Cloud | <https://cloud.google.com/blog/products/gcp/introducing-google-cloud-iot-core-for-securely-connecting-and-managing-iot-devices-at-scale>, May 2017, No author. cited by applicant

“IoT at the Edge: Bringing intelligence to the edge using Cloud IoT (Cloud Next '18)”, YouTube; Google Cloud Tech <https://www.youtube.com/watch?v=-T9MNR-BI81>, Aug. 17, 2018, pp. 1, No author. cited by applicant

“IoT Core device-to-device communication”, Google |<https://cloud.google.com/community/tutorials/iot-device-to-device>, Dec. 2017, No author. cited by applicant

“IoT Partner Quickstart”, https://services.google.com/fh/files/misc/iot_partner_quickstart1.0.pdf, Nov. 15, 2018, pp. 1-13, No author. cited by applicant

“IoT: Understanding the shift from cloud to edge computing”, <https://internetofbusiness.com/shift-from-cloud-to-edge-computing/>, Aug. 14, 2017, pp. 1-9, No author. cited by applicant

“Journey to OpenShift in a Multi-Cloud Environment, Part 3”, RedHat Hybrid Cloud | Blog <https://cloud.redhat.com/blog/journey-openshift-multi-cloud-environment-part-3>, Jan. 2, 2018, pp. 1-4, No author. cited by applicant

“Meeting at the Edge with VMware Internet of Things (IoT)”, <https://blogs.vmware.com/edge/files/2017/04/Meeting-at-the-Edge.pdf>, Apr. 2017, pp. 1-9, No author. cited by applicant

“Microsoft Azure IoT Reference Architecture”, Version 2.0 Release, May 2, 2018, pp. 1-79, No author. cited by applicant

“Microsoft Azure IoT Reference Architecture”, https://download.microsoft.com/download/A/4/D/A4DAD253-BC21-41D3-B9D9-87D2AE6F0719/Microsoft_Azure_IoT_Reference_Architecture.pdf, Sep. 26, 2018, pp. 18, No author. cited by applicant

“Microsoft IoT Central delivers low-code wayto build IoT solutions fast”, Microsoft | <https://azure.microsoft.com/en-us/blog/microsoft-iot-central-delivers-low-code-way-to-build-iot-solutions-fast/>, Dec. 2017, No author. cited by applicant

“Mobile Content Management-Datasheet”, <https://www.vmware.com/content/dam/digitalmarketing/vmware/en/pdf/workspace-one/vmware-airwatch-mcm-datasheet.pdf>, Mar. 12, 2019, pp. 4, No author. cited by applicant

“mPRM: An Overview”, ProSyst http://documentation.bosch-si.com/iot/PRM/v6.0/welcome/mprm_functional.html, Jun. 13, 2017, pp. 1-3, No author. cited by applicant

“OpenShift container cluster monitoring, logging, and Telemetry”, RedHat https://access.redhat.com/documentation/en-us/openshift_container_platform/4.2/html/container-native_virtualization/container-native-virtualization-user-s-guide#cnv-openshift-cluster-monitoring, Oct. 16, 2019, pp. 1-7, No author. cited by applicant

“Pivotal CF 1.3 App Dashboard and Usage Reports”, YouTube <https://www.youtube.com/watch?v=jpGUbjh8IUy>, Sep. 25, 2014, pp. 1. No author. cited by applicant

“Pivotal Cloud Foundry Documentation”, Pivotal Version 2.0 <https://resources.docs.pivotal.io/pdfs/pcf-docs-2.0.pdf>, Dec. 20, 2018, pp. 1-1591, No author. cited by applicant

“Pivotal Container Service (PKS)”, Pivotal Version 1.2 <https://resources.docs.pivotal.io/pdfs/pks-1-2.pdf>, Sep. 3, 2019, pp. 1-323, No author. cited by applicant

“Pivotal Container Service Overview”, Slideshare <https://www.slideshare.net/Pivotal/pivotal-container-service-overview>, Jul. 16, 2019, pp. 1-78, No author. cited by applicant

“Release Notes for VMware Unified Access Gateway 3.1 and 3.1.1”, https://docs.vmware.com/en/Unified-Access-Gateway/3.1/r/unified_access_gateway-31-release-notes.html, Oct. 27, 2017, pp. 1-4, No author. cited by applicant

“Runtime System”, Wikipedia, 2018, No author. cited by applicant

“Setting up ServiceNow Cloud Management”, DxSherpa <https://dxsherpa.com/blogs/setting-up-servicenow-cloud-management/>, Jul. 19, 2018, pp. 1-10, No author. cited by applicant

>“Setup Azure IoT Edge CI/CD Pipeline with ARM Agent”, <https://devblogs.microsoft.com/iotdev/setup-azure-iot-edge-ci-cd-pipeline-with-arm-agent/>, Nov. 13, 2018, pp. 1-11, No author. cited by applicant

“The Developer's Guide to Azure”, https://download.microsoft.com/download/2/C/F/2CF7401A-B9D7-4828-917D-199E0896BFE5/Azure_Developer_Guide_eBook.pdf, May 2019, pp. 64, No author. cited by applicant

“The Service Mesh Era: Architecting, Securing and Managing Microservices with Istio”, Google Cloud https://services.google.com/fh/files/misc/the_service_mesh_era_architecting_securing_and_managing_microservices_with_istio_white_paper.pdf, Feb. 5, 2019, pp. 1-40, No author. cited by applicant

“The Service Mesh Era: Architecting, Securing and Managing Microservices with Istio”, Google Cloud https://services.google.com/fh/files/misc/the_service_mesh_era_architecting_securing_and_managing_microservices_with_istio_white_paper.pdf, Mar. 6, 2019, pp. 6-40, No author. cited by applicant

“The Service Mesh Era: Architecting, Securing and Managing Microservices with Istio”, Google Cloud https://services.google.com/fh/files/misc/the_service_mesh_era_architecting_securing_and_managing_microservices_with_istio_white_paper.pdf, Jan. 22, 2019, pp. 1-40, No author. cited by applicant

“VMware Announces VMware Tanzu Portfolio to Transform the Way Enterprises Build, Run and Manage Software on Kubernetes”, VMware News & Stories <https://news.vmware.com/releases/vmware-announces-vmware-tanzu-Portfolio-to-transform-the-way-enterprises-build-run-and-manage-software-on-kubernetes>, Aug. 26, 2019, pp. 1-11, No author. cited by applicant

“VMware Enterprise PKS Architecture Overview”, YouTube <https://www.youtube.com/watch?v=504FGHukY8Y>, Feb. 9, 2018, pp. 1, No author. cited by applicant

“VMware Identity Manager API”, <https://developer.vmware.com/apis/57/#api>, Aug. 2018, pp. 2, No author. cited by applicant

“VMware Tanzu Mission Control Demo”, YouTube <https://www.youtube.com/watch?v=7m9S4HilJlo>, Aug. 28, 2019, pp. 1, No author. cited by applicant

“VMware Tunnel (Per-App VPN) Connections in Cascade Mode”, <https://docs.vmware.com/en/Unified-Access-Gateway/3.10/com.vmware.uag.310-deploy-config.doc/GUID-B69DD1E7-ED88-4574-912C-DBA071BC9D6F.html>, May 31, 2019, pp. 2, No author. cited by applicant

“VMware Unified Access Gateway: Use Cases—Feature Walk-through”, <https://www.youtube.com/watch?v=xan9mYOJqAs>, Oct. 27, 2017, pp. 1, No author. cited by applicant

“What is AWS IoT?”, AWS | Youtube: https://www.youtube.com/watch?v=WAp6FHbhYCK&ab_channel=AmazonWebServices ; Timestamp 4:55/10:08, Jan. 2018, No author. cited by applicant

“What is IoT Edge?”, Code Project | <https://www.codeproject.com/Articles/1261285/What-is-IoT-Edge>, Sep. 2018, pp. 1-9, No author. cited by applicant

“Workspace ONE UEM 1810 introduces support for Android Enterprise fully managed devices with work profiles”, <https://bayton.org/blog/2018/10/workspace-one-uem-1810-introduces-Support-for-android-enterprise-fully-managed-devices-with-work-profiles>, Oct. 29, 2018, pp. 11, No author. cited by applicant

Angelas , “Java Heap Space vs. Stack Memory: How Java Applications Allocate Memory”, stackify.com, Sep. 5, 2017, pp. 1-3. cited by applicant

Abadi, “TensorFlow: Large-Scale Machine Learning on Heterogeneous Distributed Systems”, Preliminary White Paper, <http://download.tensorflow.org/paper/whitepaper2015.pdf>, Nov. 9, 2015, pp. 1-19. cited by applicant

Al, Yuan et al., “Edge computing technologies for Internet of Things: a primer”, Digital Communications and Networks 4 | <https://doi.org/10.1016/j.dcan.2017.07.001>, 2018, pp. 77-86. cited by applicant

Asanghanwa, Eustace , “Simplifying confidential computing: Azure IoT Edge security with enclaves—Public preview”, <https://azure.microsoft.com/en-us/blog/simplifying-confidential-computing-azure-iot-edge-security-with-enclaves-public-preview/>, Nov. 19, 2018, pp. 1-7. cited by applicant

Avram, Abel , “AWS Greengrass Runs Lambda Functions on IoT Devices”, InfoQ | <https://www.infoq.com/news/2017/06/aws-greengrass/>, Jun. 2017, pp. 1-2. cited by applicant

Banka, Roank , “Pivotal CloudFoundry on Google cloud platform”, Slideshare <https://www.slideshare.net/RonakBanka/pcfcpgoogleioextended>, Jul. 16, 2018, pp. 1-44. cited by applicant

Barr, Jeff , “AWS Greengrass—Run AWS Lambda Functions on Connected Devices”, <https://aws.amazon.com/blogs/aws/category/aws-greengrass/>, Jun. 7, 2017, pp. 1-18. cited by applicant

Bayton, Jason , “Android Enterprise COPE enrolment comparison: WS1 UEM / MI Core”, <https://www.youtube.com/watch?v=QsJu3xFzjMw>, Oct. 28, 2018, pp. 1. cited by applicant

Beguelin, “Turn Your Smartphone Into an IoT Device”, “IBM Developer <https://developer.ibm.com/tutorials/iot-mobile-phone-iot-device-bluemix-apps-trs/>”, Nov. 10, 2015, pp. 1-12. cited by applicant

Beltre, Angel et al., “Enabling HPC workloads on Cloud Infrastructure using Kubernetes Container Orchestration Mechanisms”, 2019 IEEE/ACM Workshop on Containers and New Orchestration Paradigms for Isolated Environments in HPC (CANOPIEHPC), 2019, pp. 11-20. cited by applicant

Ben-David, Jacob , “Google Cloud’s Anthos—Everything You Need To Know”, Turbonomic Blog <https://blog.turbonomic.com/google-clouds-anthos>, Apr. 15, 2019, pp. 1-9. cited by applicant

Benson, Mark , “Technical Introduction to VMware Unified Access Gateway for Horizon Secure Remote Access”, VMware | <https://blogs.vmware.com/euc/2015/09/what-is-vmware-unified-access-gateway-secure-remote-access.html>, Sep. 9, 2015, pp. 8. cited by applicant

Berdy, Nicole , “Azure IoT Hub message routing dramatically simplifies IoT solution development”, <https://azure.microsoft.com/en-us/blog/azure-iot-hub-message-routing-enhances-device-telemetry-and-optimizes-iot-infrastructure-resources/>, Dec. 14, 2016. cited by applicant

Cano, Ignacio et al., “Curator: Self-Managing Storage for Enterprise Clusters”, University of Washington, Mar. 2017, pp. all. cited by applicant

Chi, Chrissie , “Enabling more device management scenarios with new features in IoT Hub”, Microsoft Azure | <https://azure.microsoft.com/en-us/blog/enabling-more-device-management-scenarios-with-new-features-in-iot-hub/>, May 7, 2018, pp. 1-6. cited by applicant

Decarlo, Paul , “Using Cognitive Services Containers with Azure IoT Edge”, <https://dev.to/azure/using-cognitive-services-containers-with-azure-iot-edge-1e5a>, May 18, 2019, pp. 1-15. cited by applicant

Garcin-Beldowski, Sébastien , “AWS Greengrass the forefront of edge computing”, <https://medium.com/smileinnovation/aws-greengrass-the-forefront-of-edge-computing-8ec2098a33b7>, Dec. 14, 2018, pp. 1-18. cited by applicant

George, Sam , “Azure IoT Edge generally available for enterprise-grade, scaled deployments”, <https://azure.microsoft.com/en-us/blog/azure-iot-edge-generally-available-for-enterprise-grade-scaled-deployments/>, Jun. 27, 2018, pp. 1-5. cited by applicant

Goldsborough, “A Tour of TensorFlow: Proseminar Data Mining”, Technische Universität München, <https://arxiv.org/pdf/1610.01178.pdf>, Oct. 2016, pp. 1-16. cited by applicant

Goodison, Donna , “Google Cloud Unleashes Managed Service Mesh, Serverless For Anthos”, The Channel Co. CRN <https://www.crn.com/news/cloud/google-cloud-unleashes-managed-service-mesh-serverless-for-anthos>, Sep. 16, 2019, pp. 1-6. cited by applicant

Grammatikou, “GEMBus As a Service Oriented Platform for Cloud-Based Composable Services”, 2011 IEEE Third International Conference on Cloud Computing Technology and Science, Nov. 1, 2011, pp. 666-671. cited by applicant

Harnik, Danny et al., “Secure Access Mechanism for Cloud Storage”, vol. 12, No. 3, pp. 317-336. [Retrieved from internet on Feb. 14, 2023], <<https://scpe.org/index.php/scpe/article/view/727>>, 2011, pp. 317-336. cited by applicant

Hermann, “Meet Michelangelo: Uber’s Machine Learning Platform”, Uber Engineering, <https://eng.uber.com/michelangelo-machine-learning-platform/>, Sep. 5, 2017, pp. 1-17. cited by applicant

Hoare, “A Semantic-Agent Framework For PaaS Interoperability”, 2016 International IEEE Conferences on Ubiquitous Intelligence & Computing, Advanced and Trusted Computing, Scalable Computing and Communications, Cloud and Big Data Computing, Internet of People, and Smart World Congress, Jul. 18, 2016, pp. 788-793. cited by applicant

Holzle, Urs et al., “Introducing Anthos: An entirely new platform for managing applications in today’s multi-cloud world”, Google Cloud Blog <https://cloud.google.com/blog/topics/hybrid-cloud/new-platform-for-managing-applications-in-todays-multi-cloud-world>, Apr. 9, 2019, pp. 1-6. cited by applicant

Holzle, Urs et al., “Introducing Anthos: An entirely new platform for managing applications in today’s multi-cloud world”, Google Cloud <https://cloud.google.com/blog/topics/hybrid-cloud/new-platform-for-managing-applications-in-todays-multi-cloud-world>, Apr. 9, 2019, pp. 1-7. cited by applicant

Islam, Tariq et al., “5 frequently asked questions about Google Cloud Anthos”, Google Cloud <https://cloud.google.com/blog/topics/hybrid-cloud/5-frequently-asked-questions-about-google-cloud-anthos>, Jun. 20, 2019, pp. 1-5. cited by applicant

Iyengar, Ashok , “Architecting at the Edge”, <https://www.ibm.com/cloud/blog/architecting-at-the-edge>, Oct. 21, 2019, pp. 1-14. cited by applicant

Iyengar, Ashok et al., “Architectural Decisions at the Edge”, <https://www.ibm.com/cloud/blog/architectural-decisions-at-the-edge>, Jul. 26, 2019, pp. 1-16. cited by applicant

Iyengar, Ashok et al., “Automation at the Edge”, <https://www.ibm.com/cloud/blog/automation-at-the-edge>, Feb. 18, 2017, pp. 1-13. cited by applicant

Iyengar, Ashok , “Cloud at the Edge”, <https://www.ibm.com/cloud/blog/cloud-at-the-edge>, Feb. 26, 2019, pp. 1-9. cited by applicant

Iyengar, Ashok , “DevOps at the Edge”, <https://www.ibm.com/cloud/blog/devops-at-the-edge>, Dec. 3, 2015, pp. 1-13. cited by applicant

Iyengar, Ashok , “GitOps at the Edge”, <https://www.ibm.com/cloud/blog/gitops-at-the-edge>, Nov. 2, 2017, pp. 1-13. cited by applicant

Iyengar, Ashok et al., “Rounding Out the Edges”, <https://www.ibm.com/cloud/blog/rounding-out-the-edges>, May 7, 2019, pp. 1-11. cited by applicant

Jain, “Dynamic Space-Time Scheduling for GPU Inference”, Massachusetts Institute of Technology, University of California, Berkeley; [http://learningsys.org/nips18/assets/papers/102CameraReadySubmissionGPU_Virtualization%20\(8\).pdf](http://learningsys.org/nips18/assets/papers/102CameraReadySubmissionGPU_Virtualization%20(8).pdf), Dec. 31, 2018, pp. 1-8. cited by applicant

Khaddar, “Smartphone: the Ultimate IoT and IoE Device”, IntechOpen, “Smartphones from an Applied Research Perspective”, <http://dx.doi.org/10.5772/intechopen.69734>, Nov. 2, 2017, pp. 137-162 (Ch. 7). cited by applicant

Lacoste, “User-Centric Security And Dependability In The Clouds-of-Clouds”, IEEE Cloud Computing, Sep. 2016, 64-75. cited by applicant

Lawton, George , “With AWS Greengrass, IoT apps become seamless edge to cloud”, <https://www.techtarget.com/iotagenda/feature/With-AWS-Greengrass-IoT-apps-become-seamless-edge-to-cloud>, Jul. 16, 2018, pp. 1-7. cited by applicant

Lewis, Sarah , “Cloudify”, TechTarget <https://www.techtarget.com/searchcloudcomputing/definition/Cloudify>, United Lex Competitive Analysis Report, Jun. 2019, pp. 1. cited by applicant

Li, “Scaling Machine Learning as a Service”, Uber Technologies, Inc., JMLR: Workshop and Conference Proceeding, <http://proceedings.mlr.press/v67/li17a/li17a.pdf>, Jun. 19-24, 2016, pp. 16-29. cited by applicant

Lin, Jennifer et al., “Anthos simplifies application modernization with managed service mesh and serverless for your hybrid cloud”, Google Cloud <https://cloud.google.com/blog/topics/hybrid-cloud/anthos-simplifies-application-modernization-with-managed-service-mesh-and-serverless-for-your-hybrid-cloud>, Sep. 16, 2019, pp. 1-6. cited by applicant

Liu, Gus , “How to Install a Face Recognition Model at the Edge with AWS IoT Greengrass”, <https://aws.amazon.com/blogs/iot/how-to-install-a-face-recognition-model-at-the-edge-with-aws-iot-greengrass/>, Apr. 30, 2019, pp. 1-10. cited by applicant

Lobo, Savia , “Microsoft Azure IoT Edge is open source and generally available!”, PacktHub | <https://hub.packtpub.com/microsoft-azure-iot-edge-is-open-source-and-generally-available/>, Jun. 29, 2018, pp. 1-2. cited by applicant

Malasi, Aman , “Google Anthos: Write Once, Run Anywhere”, HCL Tech Blogs <https://www.hcltech.com/blogs/google-anthos-write-once-run-anywhere>, Aug. 26, 2019, pp. 1-2. cited by applicant

Marko, Kurt , “SAP to the fore as Cloud Foundry grows into the preferred platform for cloud-native enterprise apps”, Diginomica

<https://diginomica.com/cloud-foundry-growing-preferred-platform-cloud-native-enterprise-applications>, Jun. 19, 2017, pp. 1-10. cited by applicant

McConville, Anton et al., “A brief history of Kubernetes, OpenShift, and IBM”, IBM Developer Blog | <https://developer.ibm.com/blogs/a-brief-history-of-red-hat-openshift/>, Aug. 1, 2019, pp. 9. cited by applicant

McLuckie, Craig , “Introducing VMware Tanzu Mission Control to Bring Order to Cluster Chaos”, VMware Tanzu <https://tanzu.vmware.com/content/blog/introducing-vmware-tanzu-mission-control-to-bring-order-to-cluster-chaos>, Aug. 26, 2019, pp. 1-6. cited by applicant

Mijumbi, et al., “Learning Algorithms for Dynamic Resource Allocation in Virtualised Networks”, Proceedings of Workshop on Management of Large Scale Virtualized Infrastructures: Smart Data Acquisition, Analysis and Network and Service Management in the Future Internet, 2014, 4 pages. cited by applicant

Mohamed, Riaz , “Deploy Cloud Agnostic Applications with VMware vRealize Automation Cloud”, VMware Cloud Management <https://blogs.vmware.com/management/2019/09/cloud-agnostic-apps-vra-cloud.html>, Sep. 17, 2019, pp. 1-10. cited by applicant

Morabito, Roberto et al., “A Framework Based on SDN and Containers for Dynamic Service Chains on IoT Gateways”, Proceedings of the Workshop on Hot Topics in Container Networking and Networked Systems. Ericsson Research, NomadicLab, <https://dl.acm.org/doi/pdf/10.1145/3094405.3094413>, Aug. 11, 2017, pp. 42-47. cited by applicant

MSV, Janakiram , “5 Reasons Why Azure IoT Edge Is Industry's Most Promising Edge Computing Platform”, Forbes | <https://www.forbes.com/sites/janakirammsv/2018/07/01/5-reasons-why-azure-iot-edge-is-industrys-most-promising-edge-computing-platform/?sh=56b9ef223249>, Jul. 2, 2018, pp. 1-5. cited by applicant

MSV, Janakiram , “Azure IoT Edge: A Technology Primer”, TheNewsStack | <https://thenewstack.io/azure-iot-edge-a-technology-primer/>, Sep. 14, 2018, pp. 1-9. cited by applicant

MSV, Janakiram , “Google Forays Into Edge Computing With Cloud IoT Edge And TPU”, Forbes <https://www.forbes.com/sites/janakirammsv/2018/07/30/google-forays-into-edge-computing-through-cloud-iot-edge-and-tpu/?sh=5e68b3fc6005>, Jul. 30, 2018, pp. 1-4. cited by applicant

O'Keefe, Megan , “Edge Computing and the Cloud-Native Ecosystem”, TheNewStack | <https://thenewstack.io/edge-computing-and-the-cloud-native-ecosystem/>, Apr. 18, 2018, pp. 1-11. cited by applicant

O'Keefe, Megan , “Welcome to the service mesh era: Introducing a new Istio blog post series”, Google Cloud <https://cloud.google.com/blog/products/networking/welcome-to-the-service-mesh-era-introducing-a-new-istio-blog-post-series>, Jan. 22, 2019, pp. 1-5. cited by applicant

Oleniczak, Kevin , “Using AWS IoT for Predictive Maintenance”, AWS | <https://aws.amazon.com/blogs/iot/using-aws-iot-for-predictive-maintenance/>, Jun. 28, 2018, pp. 1-6. cited by applicant

Paladi, Nicolae et al., “Domain Based Storage Protection with Secure Access Control for the Cloud”, <https://dl.acm.org/doi/pdf/10.1145/2600075.2600082>, 2014, pp. 35-42. cited by applicant

Param, Sunil , “Google's Coral: A new product development platform with local AI”, TechGig | <https://content.techgig.com/technology/googles-coral-a-new-product-development-platform-with-local-ai/articleshow/69042955.cms>, Apr. 26, 2019, pp. 1-18. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <https://nutanixbible.com/>, Apr. 9, 2019, pp. all. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <http://stevenpoitras.com/the-nutanix-bible/> (Publication date based on indicated capture date by Archive.org; first publication date unknown), Jan. 11, 2014, pp. all. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <https://nutanixbible.com/>, Jan. 12, 2016, pp. all. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <https://nutanixbible.com/>, Jan. 3, 2017, pp. all. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <https://nutanixbible.com/>, Jan. 3, 2018, pp. all. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <http://stevenpoitras.com/the-nutanix-bible/> (Publication date based on indicated capture date by Archive.org; first publication date unknown), Jan. 7, 2015, pp. all. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <https://nutanixbible.com/>, Jan. 8, 2019, pp. all. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <https://nutanixbible.com/>, Jul. 25, 2019, pp. all. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <http://stevenpoitras.com/the-nutanix-bible/> (Publication date based on indicated capture date by Archive.org; first publication date unknown), Jun. 20, 2014, pp. all. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <https://nutanixbible.com/>, Jun. 25, 2018, pp. all. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <https://nutanixbible.com/>, Jun. 8, 2017, pp. all. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <http://stevenpoitras.com/the-nutanix-bible/> (Publication date based on indicated capture date by Archive.org; first publication date unknown), Jun. 9, 2015, pp. all. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <https://nutanixbible.com/>, Jun. 9, 2016, pp. all. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <http://stevenpoitras.com/the-nutanix-bible/> (Publication date based on indicated capture date by Archive.org; first publication date unknown), Oct. 15, 2013, pp. all. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <https://nutanixbible.com/>, Sep. 17, 2019, pp. all. cited by applicant

Poitras, Steven , “The Nutanix Bible”, <https://nutanixbible.com/>, Sep. 4, 2015, pp. all. cited by applicant

Rafique, “Towards Portability and Interoperability Support in Middleware for Hybrid Clouds”, 2014 IEEE Infocom Workshop on Cross-Cloud Systems, Apr. 27, 2014, pp. 7-12. cited by applicant

Rai, Rahul , “Building Applications with Azure IoT Edge”, <https://thecloudblog.net/post/building-applications-with-azure-iot-edge/>, Apr. 25, 2018, pp. 1-16. cited by applicant

Rai, Rahul , “IoT Edge Device Monitoring and Management with Azure Durable Entities Functions—Part 1”, <https://thecloudblog.net/post/iot-edge-device-monitoring-and-management-with-azure-durable-entities-functions-part-1/>, Jul. 1, 2019, pp. 1-14. cited by applicant

Reid, Nate , “Upgrading VMware Enterprise PKS 1.2 to 1.3”, YouTube <https://www.youtube.com/watch?v=7pAxNWIXVLM>, Jan. 28, 2019, pp. 1. cited by applicant

Ren, Ju et al., “Edge Computing for the Internet of Things”, IEEE Journals & Magazine | vol. 32 Issue: 1, 2008, pp. 1-6. cited by applicant

Rhee, Injong , “Bringing intelligence to the edge with Cloud IoT”, Google Cloud | <https://cloud.google.com/blog/products/gcp/bringing-intelligence-edge-cloud-iot>, Jul. 25, 2018, pp. 1-7. cited by applicant

Rosoff, Jared , “Project Pacific—Technical Overview”, VMware | <https://blogs.vmware.com/vsphere/2019/08/project-pacific-technical-overview.html>, Aug. 26, 2019, pp. 7. cited by applicant

Schonbaum, Iftach , “Anthos—Google's Enterprise Hybrid & Multi-Cloud Platform”, CloudZone <https://medium.com/cloudzone/anthos-google>

enterprise-hybrid-multi-cloud-platform-7975e05a7729, Apr. 23, 2019, pp. 1-7. cited by applicant

Stakun, Jaroslaw , “Introduction to Red Hat OpenShift Container Platform”, Open Virtualization Pro <https://www.openvirtualization.pro/red-hat-openshift-container-platform/>, Apr. 26, 2019, pp. 1. cited by applicant

Stoks, “Workspace IoT Series: How Industry 4.0 Will Transform Your Mobile Strategy”, “VMware End-User Computing Blog <https://blogs.vmware.com/euc/2019/02/workspace-iot-mobile-strategy.html>”, Feb. 20, 2019, pp. 1-5. cited by applicant

Sun, “Towards Distributed Machine Learning in Shared Clusters: a Dynamically-Partitioned Approach”, Cornell University: Arxiv: 1704.06738v1 [CS.CD], Apr. 22, 2017, 6 pages. cited by applicant

Tamura, Yoshi , “GPUs as a service with Kubernetes Engine are now generally available”, Google Cloud | <https://cloud.google.com/blog/products/gcp/gpus-service-kubernetes-engine-are-now-generally-available>, Jun. 19, 2018, pp. 5. cited by applicant

Vanderzyden, John , “Using AWS Greengrass to Enable IoT Edge Computing”, mabl | <https://www.mabl.com/blog/using-aws-greengrass-enable-iot-edge-computing>, Aug. 23, 2017, pp. 1-12. cited by applicant

Warburton, “An Intro to GPU Architecture and Programming Models I Tim Warburton, Virginia Tech”, YouTube, Argonne National Laboratory Training, <https://www.youtube.com/watch?v=IGmPy8xpT4E>, Sep. 25, 2017, pp. 1-3. cited by applicant

Warke, Amit et al., “Storage Service Orchestration with Container Elasticity”, 2018 IEEE 4th International Conference on Collaboration and Internet Computing DOI 10.1109/CIC.2018.00046, 2018, pp. 283-292. cited by applicant

Wiggers, Kyle , “Microsoft launches Azure IoT Edge out of preview”, VentureBeat | <https://venturebeat.com/mobile/microsoft-launches-azure-iot-edge-out-of-preview/>, Jun. 27, 2018, pp. 1-6. cited by applicant

Wiggers, Steef-Jan , “Google Cloud Platform Release Edge TPU and Cloud IoT Edge”, InfoQ <https://www.infoq.com/news/2018/07/google-iot-edge-tpu-ai/>, Jul. 30, 2018, pp. 1-4. cited by applicant

Wiggers, Steef-Jan , “Google Kubernetes Engine 1.10 Is Generally Available and Enterprise Ready”, InfoQ | <https://www.infoq.com/news/2018/06/google-kubernetes-engine-1.10-ga>, Jun. 1, 2018, pp. 4. cited by applicant

Yamashita, Teppei , “Real-time data processing with IoT Core”, Google Cloud | <https://cloud.google.com/community/tutorials/cloud-iot-rtdp>, Apr. 2018, pp. 1-19. cited by applicant

Zhang, Xinyi , “Create a CI/CD pipeline for your IoT Edge solution with Azure DevOps”, Microsoft | <https://devblogs.microsoft.com/iotdev/create-a-ci-cd-pipeline-for-your-iot-edge-solution-with-azure-devops/>, Oct. 29, 2018, pp. 1-10. cited by applicant

Zhao, Zhuoran et al., “DeepThings: Distributed Adaptive Deep Learning Inference on Resource-Constrained IoT Edge Clusters”, IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems, vol. 37, No. 11 | doi: 10.1109/TCAD.2018.2858384., Nov. 2018, pp. 2348-2359. cited by applicant

Primary Examiner: Ho; Andy

Attorney, Agent or Firm: Stoel Rives LLP

Background/Summary

CROSS REFERENCE TO RELATED APPLICATION(S) (1) This application is a continuation of U.S. application Ser. No. 18/185,300 filed Mar. 16, 2023, which is a divisional of U.S. application Ser. No. 16/666,242 filed Oct. 28, 2019, issued as U.S. Pat. No. 11,635,990 on Apr. 25, 2023, which claims priority to provisional application No. 62/869,348, filed Jul. 1, 2019, which are incorporated herein by reference, in their entirety, for any purpose.

BACKGROUND

(1) Internet of Things (IoT) systems are increasing in popularity. Generally, IoT systems utilize a number of edge devices. Edge devices may generally refer to computing systems deployed about an environment (which may be a wide geographic area in some examples). The edge devices may include computers, servers, clusters, sensors, appliances, vehicles, communication devices, etc. Edge devices may obtain data (including sensor data, voice data, image data, and/or video data, etc.). While edge devices may provide some processing of the data at the edge device, in some examples edge devices may be connected to a centralized analytics system (e.g., in a cloud or other hosted environment). The centralized analytics system, which may itself be implemented by one or more computing systems, may further process data received from edge devices by processing data received by individual edge devices and/or by processing combinations of data received from multiple edge devices.

(2) Within an IoT computing environment, real-time data analytics computations at edge systems of an IoT system can be complex due to the distributed nature of edge system computing and diversity in methods and protocols for ingesting, processing, and communicating information back to a data cloud. That is, edge systems may be deployed over a wide geographic area, and diversity among the edge systems may make deployment of applications to retrieve and consume edge data across multiple edge systems of the IoT system difficult.

(3) As the number of and complexity of edge systems advantageously increases, a workload of a control plane to manage the IoT system may also increase to a point where compute resources available to the control plane are insufficient to effectively manage the IoT system. When the compute resources become insufficient, the reliability of the IoT system may be adversely affected.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is a block diagram of an Internet of Things system, in accordance with an embodiment of the present disclosure.
- (2) FIG. 2 is a block diagram of a distributed computing system, in accordance with an embodiment of the present disclosure.
- (3) FIG. 3 is a block diagram of an IoT system topology, in accordance with an embodiment of the present disclosure.
- (4) FIG. 4A is a flow diagram of an example process that may be implemented to provide software update in an edge device in accordance with an embodiment of the present disclosure.

(5) FIG. 4B is a flow diagram of an example process that may be implemented in an edge device to establish a secure communication between an edge device and an IoT manager in accordance with an embodiment of the present disclosure.

(6) FIG. 5 is a block diagram of a centralized IoT manager of an IoT system, in accordance with an embodiment of the present disclosure.

(7) FIG. 6 is a block diagram illustrating communication via a centralized IoT manager of an IoT system, in accordance with an embodiment of the present disclosure.

(8) FIG. 7 is a flow diagram of a method to communicate with an edge system of an IoT system, in accordance with an embodiment of the present disclosure.

(9) FIG. 8 is a flow diagram of a method to provide a software update to an edge system of an IoT system, in accordance with an embodiment of the present disclosure.

(10) FIG. 9 is a block diagram of components of an edge system or computing in accordance with an embodiment of the present disclosure.

DETAILED DESCRIPTION

(11) Examples described herein include a scalable centralized Internet of Things (IoT) manager of an IoT system capable of expanding as a number of tenants and edge systems within an IoT system increases. The centralized IoT manager may be implemented using a software as a service (SaaS) architecture, in some examples. In some examples, the centralized IoT manager may support two-way communication using secure protocols, such as hyper-text transfer protocol secure (HTTPS), WebSocket secure (WSS), etc. A metadata object model may be provided to fully specify behavior of edge systems. The centralized IoT manager may manage an initial synchronization of the object model of the edge system, which may then persist at the edge system and allow the edge system to continue to operate offline. Subsequent object model synchronizations may be managed using delta-based communication and update notifications, which may be more efficient than full object model synchronizations. In some examples, an edge system may include a single edge device, an edge virtual machine (VM) hosted on a server or other computing device, or an edge cluster (e.g., tenant) that includes a cluster of edge devices.

(12) The centralized IoT manager may include a stateless service that is hosted on a computing node or a computing node cluster. The centralized IoT manager may include multiple instances of one or more services to facilitate scalability of the IoT system. The centralized IoT manager is configured to centrally manage edge systems of the IoT system, including pushing software and/or configuration updates to the edge systems, providing data pipelines and other applications to the edge systems, storing/updating respective metadata for the IoT system in a database, etc.

(13) The centralized IoT manager may communicate with the edge systems using a secure, two-way communication protocol, such as HTTPS, WSS, etc. For example, the centralized IoT manager may open a respective secure socket channel to communicate with each of the edge devices. The centralized IoT manager may use representational state transfer (REST) application programming interfaces (API) or remote procedure calls (RPCs) to communicate with the edge systems. In addition, rather than providing a full software package to each of the edge systems for an update to an application hosted on an edge system, the centralized IoT manager may provide delta changes between the new version and a current version of the application to each edge system to make updating edge systems more efficient. The centralized IoT manager may manage the software updates of the edge systems by pushing the software updates according to a schedule managed by the centralized IoT manager, rather than being managed by the edge systems periodically checking for updates.

(14) Various embodiments of the present disclosure will be explained below in detail with reference to the accompanying drawings. The detailed description includes sufficient detail to enable those skilled in the art to practice the embodiments of the disclosure. Other embodiments may be utilized, and structural, logical and electrical changes may be made without departing from the scope of the present disclosure. The various embodiments disclosed herein are not necessary mutually exclusive, as some disclosed embodiments can be combined with one or more other disclosed embodiments to form new embodiments.

(15) FIG. 1 is a block diagram of an Internet of Things (IoT) system **100**, in accordance with an embodiment of the present disclosure. The IoT system **100** may include one or more of any of edge cluster(s) **110** coupled to respective data source(s) **120**, edge device(s) **112** coupled to respective data source(s) **122**, a server/cluster **114** coupled to respective data source(s) **124** and configured to host one or more edge virtual machines VM(s) **115**. In some examples, an edge system may include a single edge device, an edge virtual machine (VM) hosted on a server or other computing device, or an edge cluster (e.g., tenant) that includes a cluster of edge devices.

(16) The IoT system **100** may further include a central IoT computing system **140** coupled to the one or more of the edge cluster(s) **110**, the edge device(s) **112**, and/or the edge VM(s) **115** hosted on the server/cluster **114** via a network **130** to manage configuration and operation of the IoT system **100**. The IoT system **100** may further include a data computing system **150** coupled to the network **130** and configured to receive, store, process, etc., data received from the one or more of the edge cluster(s) **110**, the edge device(s) **112**, and/or the server/cluster **114** via a network **130**.

(17) The network **130** may include any type of network capable of routing data transmissions from one network device (e.g., the edge cluster(s) **110**, the edge device(s) **112**, the server/cluster **114**, a computing node of the central IoT computing system **140**, and/or a computing node of the data computing system **150**) to another. For example, the network **130** may include a local area network (LAN), wide area network (WAN), intranet, or a combination thereof. The network **130** may include a wired network, a wireless network, or a combination thereof.

(18) The IoT system **100** may include one or more types of edge systems selected from any combination of the edge cluster(s) **110**, the edge device(s) **112**, and/or the edge VM(s) **115** hosted on the server/cluster **114**. Each of the edge cluster(s) **110** (e.g., tenants) may include a respective cluster of edge nodes or devices that are configured to host a respective edge stack **111**. The edge stack **111** may be distributed across multiple edge nodes, devices, or VMs of a respective one of the edge cluster(s) **110**, in some examples. Each of the edge device(s) **112** may be configured to host a respective edge stack **113**. Each of the edge VM(s) **115** may be configured to host a respective edge stack **116**. In some examples, the server/cluster **114** may be included as part of the central IoT computing system **140** or the data computing system **150**. For clarity, “edge” or “edge system” may refer to any of the edge cluster(s) **110**, the edge device(s) **112**, and/or the edge VM(s) **115** hosted on the server/cluster **114**. The edge stacks (e.g., any of the edge stack **111**, the edge stack **113**, and/or the edge stack **116**) may include software configured to operate the respective edge or edge system in communication between one or more of the respective data sources (e.g., the data source(s) **120**, the data source(s) **122**, and/or the data source(s) **124**). The software may include instructions that are stored on a computer readable medium (e.g., memory, disks, etc. that is executable by one or more processor units (e.g., central processor units (CPUs), graphic processor units (GPUs), tensor processing units (TPUs), hardware

accelerators, video processing units (VPUs), etc.) to perform functions, methods, etc., described herein.

(19) The data source(s) **120**, the data source(s) **122**, and the data source(s) **124** (“data sources”) may each include one or more devices configured to receive and/or generate respective source data. The data sources may include sensors (e.g., electrical, temperature, matter flow, movement, position, biometric data, or any other type of sensor), cameras, transducers, any type of RF receiver, or any other type of device configured to receive and/or generate source data.

(20) Each of the edge stacks may include one or more data pipelines and/or applications. In some examples, some data pipelines and/or applications may be configured to receive and process/transform source data from one or more of the data sources, other data pipelines, or combinations thereof. In some examples, a data pipeline may span across multiple edge systems. Each of the one or more data pipelines and/or applications may be configured to process respective received data based on respective algorithms or functions to provide transformed data. The data pipelines can be constructed using computing primitives and building blocks, such as VMs, containers, processes, or any combination thereof. In some examples, the data pipelines may be constructed using a group of containers (e.g., a pod) that each perform various functions within the data pipeline (e.g., subscriber, data processor, publisher, connectors that transform data for consumption by another container within the application or pod, etc.) to consume, transform, and produce messages or data. In some examples, the definition of stages of a constructed data pipeline application may be described using a user interface or API, with data ingestion and movement handled by connector components built into the data pipeline. Thus, data may be passed between containers of a data pipeline using API calls. In some examples, the edge stacks may further include respective machine learning (ML) inference services that are configured to load and execute respective ML model inferences to provide inference data that may be stored or provided at an output.

(21) The centralized IoT manager **142** hosted on the central IoT computing system **140** may be configured to centrally manage configuration of each of the edge systems and data sources via a central control plane. The central IoT computing system **140** may include one or more computing nodes configured to host the centralized IoT manager **142**. In some examples, the centralized IoT manager **142** may be distributed across a cluster of computing nodes of the central IoT computing system **140**.

(22) In some examples, the centralized IoT manager **142** may be configured to manage, for each of the edge systems, network configuration and security protocols, installed software (e.g., including data pipelines and applications), connected data source(s) (e.g., including type, category, identifiers, data communication protocols, etc.), connected data plane(s), communication between the edge systems and users, etc. The centralized IoT manager **142** may maintain configuration information for each of the edge systems, data sources, associated users, including hardware configuration information, installed software version information, connected data source information (e.g., including type, category, identifier, etc.), associated data planes, current operational status, authentication credentials and/or keys, etc.

(23) The centralized IoT manager **142** may be configured to generate (e.g., build, construct, update, etc.) and distribute data pipelines and applications to selected edge systems based on the configuration maintained for each edge system. For example, in response to a request for a new data pipeline associated with a particular type or category of data sources, the centralized IoT manager **142** may identify data sources having the particular type or category (e.g., or attribute), and identify respective edge systems are connected to the identified data sources of the particular type or category. For each identified edge system, the centralized IoT manager **142** may generate a respective version of the application or data pipeline based on respective hardware configuration information for the edge system. That is, the centralized IoT manager **142** may independently generate the applications and data pipelines to efficiently operate according to the specific hardware configuration of each edge system.

(24) Edge data and/or ML inference data may be provided from the edge systems to one or more respective data planes, such as the data plane **152** of the data cloud computing system **150**, users, or other edge systems via the network **130**. In some examples, the edge data may include some or all of the source data from one or more of the data sources, processed source data, data derived from the source data, combined source data, or any combination thereof. In some examples, the edge data may include and/or may be based on ML inference data. The data plane **152** may be configured to store the edge data, process the edge data, provide access to the edge data to clients, etc. The data computing system **150** may include one or more cloud platforms that includes a plurality of computing nodes configured to host one or more versions of the data plane **152**.

(25) In operation, the IoT system **100** may include any number and combination of data sources selected from the data source(s) **120**, the data source(s) **122**, and the data source(s) **124** that are each configured to provide respective source data. The data sources of the IoT system **100** may collectively span any type of geographic area (e.g., across continents, countries, states, cities, counties, facilities, buildings, floors, rooms, systems, units, or any combination thereof). The number of data sources may range in the tens, hundreds, thousands, or more. The data sources may include sensors (e.g., electrical, temperature, matter flow, movement, position, biometric data, or any other type of sensor), cameras, transducers, any type of radio frequency (RF) receiver, or any other type of device configured to receive and/or generate source data.

(26) Rather than each of the data sources independently sending all source data directly to a data plane or user, the IoT system **100** may include any number and combination of edge systems selected from any combination of the edge cluster(s) **110**, the edge device(s) **112**, and/or the edge VM(s) **115** hosted on the server/cluster **114** that are proximately located with and connected to respective data sources and are each configured to receive and select/process/transform the source data that is provided to the data plane or user. The edge systems within the IoT system **100** may include homogenous hardware and software architectures, in some examples. In other examples, the edge systems have a wide array of hardware and software architectures and capabilities. Each of the edge systems may be connected to a respective subset of data sources, and may host respective data pipelines and applications (e.g., included in the edge stacks, such as the edge stack **111**, edge stack **113**, or edge stack **116**) that are configured to process source data from a respective one or more of the connected data sources and/or transformed data from other applications and/or data pipelines.

(27) Each of the one or more data pipelines and/or applications may be configured to process and/or distribute respective transformed data based on received source data (e.g., or other edge data) using respective algorithms or functions. In some examples, the algorithms or functions may include any other user-specified or defined function to process/transform/select/etc. received data. In some examples, an edge system may provide the transformed data from a data pipeline or an application of the one or more data pipelines or applications of the edge stacks to a respective destination data plane, such as the data plane **152** of the data computing system **150** as edge data. In some examples, the edge systems may be configured to share edge data with other edge systems. The one or more data pipelines or applications of the edge stacks may be implemented using a containerized architecture that is managed via a container orchestrator. The data pipelines and/or applications communicate using API calls, in some examples. In some examples, ML inference services may be configured to receive a request for a ML model inference, and to load the requested ML model inference in an inference engine in

response to the request to execute the ML model inference to provide inference data.

(28) In some examples, the edge systems may cause transformed data from a data pipeline or an application, and/or inference data from the ML inference services to be provided to a respective data plane as edge data, such as the data plane **152** of the data computing system **150**, using respective data plane communication interfaces, including application programming interfaces (APIs). The data computing system **150** may be a dedicated computing system, or may include a centralized analytics system hosted on a network of remote servers that are configured to store, manage, and process data (e.g., cloud computing system). The centralized IoT manager **142** hosted on the central IoT computing system **140** may be configured to centrally manage configuration of each of the edge systems and data sources. In some examples, the centralized IoT manager **142** may be configured to manage, for each of the edge systems, data sources, and/or users, network configuration and security protocols, installed software (e.g., including data pipelines and applications), connected data source(s) (e.g., including type, category, identifiers, data communication protocols, etc.), connected data plane(s), etc. The centralized IoT manager **142** may maintain configuration information for each of the edge systems, data sources, associated users, including hardware configuration information, installed software version information, connected data source information (e.g., including type, category, identifier, etc.), associated data planes, current operational status, authentication credentials and/or keys, etc.

(29) The centralized IoT manager **142** may be configured to generate or update and distribute data pipelines and applications to selected edge systems based on the configuration maintained for each edge system. For example, in response to a request for a new data pipeline or application associated with a particular type or category of data sources, the centralized IoT manager **142** may identify data sources having the particular type or category, and identify respective edge systems are connected to the identified data sources of the particular type or category. For each identified edge system, the centralized IoT manager **142** may generate a respective version of the application or data pipeline based on respective hardware configuration information for the edge system. That is, the centralized IoT manager **142** may independently generate the applications and data pipelines to efficiently operate according to the specific hardware configuration of each edge system. The data pipelines may be constructed using a group of containers (e.g., a pod) each configured to perform various functions within the data pipeline (e.g., subscriber, data processor, publisher, connectors that transform data for consumption by another container within the application or pod, etc.). In some examples, the centralized IoT manager **142** may be configured to define stages of a constructed data pipeline application using a user interface or REST API, with data ingestion and movement handled by the connector components built into the data pipeline.

(30) The edge systems may provide the edge data and/or the ML inference data to one or more respective data planes, such as the data plane **152** of the data computing system **150**, via the network **130**. In some examples, the edge stacks may be configured to implement respective data plane communication interfaces, including APIs, to communicate with the one or more data planes. The data plane **152** may be configured to store the edge data, process the edge data, aggregate the edge data across the IoT system **100**, provide access to the edge data to clients, or any combination thereof. The edge data received and processed at the data plane **152** may provide information about events, trends, health, etc., of the IoT system **100** based in data captured by the data sources.

(31) In some examples, the centralized IoT manager **142** may be deployed in multiple instances that are each instance respectively configured to connect to a respective group of one or more edge systems. The centralized IoT manager **142** may further include a containerized system configured to: receive from a first instance of the plurality of instances of the centralized IoT manager **142** a request to publish a message on a communication channel; and in response to the request to publish the message, send the message to the communication channel. The communication channel may be established between a second instance of the plurality of instances of the centralized IoT manager **142** and a second edge system of the respective group of one or more edge systems. Publishing the message on the communication channel causes the second instance of the plurality of instances of the centralized IoT manager **142** to send the message to the second edge system.

(32) The containerized system of the centralized IoT manager **142** may be further configured to, upon establishment of the communication channel between the second instance of the plurality of instances of the centralized IoT manager **142** and the second edge system, receive a channel number of the communication channel from the second instance of the centralized IoT manager **142**. The centralized IoT manager **142** may receive a look up request from the first instance of the plurality of instances of the centralized IoT manager **142** for the communication channel number associated with the second edge system, and in response to the look up request, may provide the channel number of the second edge system to the first instance of the plurality of instances of the centralized IoT manager **142**. The look up request may include an identity of the second edge system associated with the channel number. In some examples, the containerized system includes a publish/subscribe-to-channel map configured to store one or more communication channel numbers subscribed for receiving update notifications. The containerized system may also include an edge-to-channel map configured to store information associating edge devices to communication channels. In providing the channel number of the second edge device to the first instance of the centralized IoT manager **142**, the containerized system may use the edge-to-channel map to determine the channel number based on the edge system.

(33) In some examples, the centralized IoT manager **142** may be configured to form secure two-way communication with an edge system of the centralized IoT manager **142**, and synchronize an object model with an edge object model for the edge system. The centralized IoT manager **142** may be further configured to maintain the edge system using delta change communications. For example, an instance of the centralized IoT manager **142** may be configured to form a communication with the edge system of the centralized IoT manager **142**, and detect a change of entity on the edge system based on an operation of the edge system. The centralized IoT manager **142**; may be further configured to determine one or more additional edge systems that have associations with the detected change of entity; and transmit a respective message to the one or more additional edge systems. Each respective message includes information about a change to a corresponding edge system of the one or more additional edge systems. The instance of the centralized IoT manager **142** may also be configured to, in response to a request for inventory from the edge system, provide the inventory of the IoT system **100** to the edge system. An edge system connected to the centralized IoT manager **142** may receive the respective message from the centralized IoT manager **142**, update an inventory of the respective edge system based on the received respective message, and synchronize an application or a data pipeline of the respective edge system with the updated inventory.

(34) In some examples, the centralized IoT manager **142** may be implemented in multiple instances. A first instance may be configured to: receive a user message to deploy a data pipeline to an edge system connected to a second instance of the centralized IoT manager **142**, determine a communication channel for the edge system, and publish a message on the determined communication channel for the edge system. A second instance may be configured to: establish a communication channel with the edge system, store a channel identification of the communication channel with association with the edge system in the centralized IoT manager **142**, and subscribe the communication channel to the centralized IoT manager **142** for receiving update notifications. The second instance may further receive

the message published by the first instance of the centralized IoT manager **142**, and in response to receiving the message, send the message to the edge system.

(35) In some examples, the centralized IoT manager **142** may receive a request from an edge system for first authentication information, provide the first authentication information to the new edge system, and receive a login request from the new edge system. The login request may contain second authentication information. The centralized IoT manager **142** may verify the second authentication information, and upon verification of the second authentication information, may establish a secure communication channel with the new edge system.

(36) FIG. 2 is a block diagram of a distributed computing system **200**, in accordance with an embodiment of the present disclosure. The distributed computing system **200** generally includes computing nodes (e.g., host machines, servers, computers, nodes, etc.) **204(1)-(N)** and storage **270** connected to a network **280**. While FIG. 2 depicts three computing nodes, the distributed computing system **200** may include two or more than three computing nodes without departing from the scope of the disclosure. The network **280** may be any type of network capable of routing data transmissions from one network device (e.g., computing nodes **204(1)-(N)** and the storage **270**) to another. For example, the network **280** may be a local area network (LAN), wide area network (WAN), intranet, Internet, or any combination thereof. The network **280** may be a wired network, a wireless network, or a combination thereof. The central IoT computing system **140** of FIG. 1 may be configured to implement the distributed computing system **200**, in some examples.

(37) The storage **270** may include respective local storage **206(1)-(N)**, cloud storage **250**, and networked storage **260**. Each of the respective local storage **206(1)-(N)** may include one or more solid state drive (SSD) devices **240(1)-(N)** and one or more hard disk drives (HDD) devices **242(1)-(N)**. Each of the respective local storage **206(1)-(N)** may be directly coupled to, included in, and/or accessible by a respective one of the computing nodes **204(1)-(N)** without communicating via the network **280**. The cloud storage **250** may include one or more storage servers that may be stored remotely to the computing nodes **204(1)-(N)** and may be accessed via the network **280**. The cloud storage **250** may generally include any type of storage device, such as HDDs, SSDs, optical drives, etc. The networked storage (or network-accessed storage) **260** may include one or more storage devices coupled to and accessed via the network **280**. The networked storage **260** may generally include any type of storage device, such as HDDs, SSDs, optical drives, etc. In various embodiments, the networked storage **260** may be a storage area network (SAN).

(38) Each of the computing nodes **204(1)-(N)** may include a computing device configured to host a respective hypervisor **210(1)-(N)**, a respective controller virtual machine (CVM) **222(1)-(N)**, respective user (or guest) virtual machines (VMs) **230(1)-(N)**, and respective containers **232(1)-(N)**. For example, each of the computing nodes **204(1)-(N)** may be or include a server computer, a laptop computer, a desktop computer, a tablet computer, a smart phone, any other type of computing device, or any combination thereof. Each of the computing nodes **204(1)-(N)** may include one or more physical computing components, such as one or more processor units, respective local memory **244(1)-(N)** (e.g., cache memory, dynamic random-access memory (DRAM), non-volatile memory (e.g., flash memory), or combinations thereof), the respective local storage **206(1)-(N)**, ports (not shown) to connect to peripheral input/output (I/O) devices (e.g., touchscreens, displays, speakers, keyboards, mice, cameras, microphones, environmental sensors, etc.).

(39) Each of the user VMs **230(1)-(N)** hosted on the respective computing node includes at least one application and everything the user VM needs to execute (e.g., run) the at least one application (e.g., system binaries, libraries, etc.). Each of the user VMs **230(1)-(N)** may generally be configured to execute any type and/or number of applications, such as those requested, specified, or desired by a user. Each of the user VMs **230(1)-(N)** further includes a respective virtualized hardware stack (e.g., virtualized network adaptors, virtual local storage, virtual memory, processor units, etc.). To manage the respective virtualized hardware stack, each of the user VMs **230(1)-(N)** is further configured to host a respective operating system (e.g., Windows®, Linux®, etc.). The respective virtualized hardware stack configured for each of the user VMs **230(1)-(N)** may be defined based on available physical resources (e.g., processor units, the local memory **244(1)-(N)**, the local storage **206(1)-(N)**, etc.). That is, physical resources associated with a computing node may be divided between (e.g., shared among) components hosted on the computing node (e.g., the hypervisor **210(1)-(N)**, the CVM **222(1)-(N)**, other user VMs **230(1)-(N)**, the containers **232(1)-(N)**, etc.), and the respective virtualized hardware stack configured for each of the user VMs **230(1)-(N)** may reflect the physical resources being allocated to the user VM. Thus, the user VMs **230(1)-(N)** may isolate an execution environment by packaging both the user space (e.g., application(s), system binaries and libraries, etc.) and the kernel and/or hardware (e.g., managed by an operating system). While FIG. 2 depicts the computing nodes **204(1)-(N)** each having multiple user VMs **230(1)-(N)**, a given computing node may host no user VMs or may host any number of user VMs.

(40) Rather than providing hardware virtualization like the user VMs **230(1)-(N)**, the respective containers **232(1)-(N)** may each provide operating system level virtualization. Thus, each of the respective containers **232(1)-(N)** is configured to isolate the user space execution environment (e.g., at least one application and everything the container needs to execute (e.g., run) the at least one application (e.g., system binaries, libraries, etc.)) without requiring an operating system to manage hardware. Individual ones of the containers **232(1)-(N)** may generally be provided to execute any type and/or number of applications, such as those requested, specified, or desired by a user. Two or more of the respective containers **232(1)-(N)** may run on a shared operating system, such as an operating system of any of the hypervisor **210(1)-(N)**, the CVM **222(1)-(N)**, or other user VMs **230(1)-(N)**. In some examples, an interface engine may be installed to communicate between a container and an underlying operating system. While FIG. 2 depicts the computing nodes **204(1)-(N)** each having multiple containers **232(1)-(N)**, a given computing node may host no containers or may host any number of containers.

(41) Each of the hypervisors **210(1)-(N)** may include any type of hypervisor. For example, each of the hypervisors **210(1)-(N)** may include an ESX, an ESX(i), a Hyper-V, a KVM, or any other type of hypervisor. Each of the hypervisors **210(1)-(N)** may manage the allocation of physical resources (e.g., physical processor units, volatile memory, the storage **270**) to respective hosted components (e.g., CVMs **222(1)-(N)**, respective user VMs **230(1)-(N)**, respective containers **232(1)-(N)**) and performs various VM and/or container related operations, such as creating new VMs and/or containers, cloning existing VMs and/or containers, etc. Each type of hypervisor may have a hypervisor-specific API through which commands to perform various operations may be communicated to the particular type of hypervisor. The commands may be formatted in a manner specified by the hypervisor-specific API for that type of hypervisor. For example, commands may utilize a syntax and/or attributes specified by the hypervisor-specific API. Collectively, the hypervisors **210(1)-(N)** may all include a common hypervisor type, may all include different hypervisor types, or may include any combination of common and different hypervisor types.

(42) The CVMs **222(1)-(N)** may provide services for the respective hypervisors **210(1)-(N)**, the respective user VMs **230(1)-(N)**, and/or the respective containers **232(1)-(N)** hosted on a respective computing node of the computing nodes **204(1)-(N)**. For example, each of the CVMs **222(1)-(N)** may execute a variety of software and/or may serve the I/O operations for the respective hypervisor **210(1)-(N)**, the respective user VMs **230(1)-(N)**, and/or the respective containers **232(1)-(N)** hosted on the respective computing node **204(1)-(N)**.

The CVMs 222(1)-(N) may communicate with one another via the network 280. By linking the CVMs 222(1)-(N) together via the network 280, a distributed network (e.g., cluster, system, etc.) of the computing nodes 204(1)-(N) may be formed. In an example, the CVMs 222(1)-(N) linked together via the network 280 may form a distributed computing environment (e.g., a distributed virtualized file server) 220 configured to manage and virtualize the storage 270. In some examples, a SCSI controller, which may manage the SSD devices 240(1)-(N) and/or the HDD devices 242(1)-(N) described herein, may be directly passed to the respective CVMs 222(1)-(N), such as by leveraging a VM-Direct Path. In the case of Hyper-V, the SSD devices 240(1)-(N) and/or the HDD devices 242(1)-(N) may be passed through to the respective CVMs 222(1)-(N).

(43) The CVMs 222(1)-(N) may coordinate execution of respective services over the network 280, and the services running on the CVMs 222(1)-(N) may utilize the local memory 244(1)-(N) to support operations. The local memory 244(1)-(N) may be shared by components hosted on the respective computing node 204(1)-(N), and use of the respective local memory 244(1)-(N) may be controlled by the respective hypervisor 210(1)-(N). Moreover, multiple instances of the same service may be running throughout the distributed system 200. That is, the same services stack may be operating on more than one of the CVMs 222(1)-(N). For example, a first instance of a service may be running on the CVM 222(1), a second instance of the service may be running on the CVM 222(2), etc.

(44) In some examples, the CVMs 222(1)-(N) may be configured to collectively manage a centralized IoT manager of an IoT system, with each of the CVMs 222(1)-(N) hosting a respective centralized IoT manager instance 224(1)-(N) on an associated operating system to form the centralized IoT manager. In some examples, one of the centralized IoT manager instances 224(1)-(N) may be designated as a master centralized IoT manager instance configured to coordinate collective operation of the centralized IoT manager instances 224(1)-(N). The centralized IoT manager instances 224(1)-(N) may be configured to manage configuration of (e.g., network connectivity information, connected data sources, installed application and other software versions, data pipelines, etc.), as well as generate and distribute data pipelines to edge systems (e.g., any of an edge device of the edge cluster(s) 110, the edge device(s) 112, the edge VM(s) 115 of the server/cluster 114, etc.) of an IoT system to centrally manage operation of the IoT system. The centralized IoT manager instances 224(1)-(N) may be configured to interface with multiple edge system types and interfaces via a control plane.

(45) To manage the operation of the IoT system, the centralized IoT manager instances 224(1)-(N) may retrieve data from and store data to IoT system data 272 of the storage 270. The IoT system data 272 may include metadata and other data corresponding to each edge system, data source, user, site, etc. within the IoT system. For example, the IoT system data 272 may include hardware configurations, software configurations, network configurations, edge system and/or data source type, categories, geographical and physical locations, authentication information, associations between edge systems and data sources, associations between edge systems and users, user access permissions, etc., or any combination thereof.

(46) In some examples, the IoT system data 272 may include metadata that is represented by an object model, which classifies metadata into different entity. Examples of entities may include a tenant (e.g., a group of one or more edge devices or systems), an edge system, a user, an application, a data source, a data pipeline, a function, a cloud profile, a container registry profile and/or a category. A tenant may include an instance per customer and tenant data key. An edge may include metadata for an edge device or system, such as serial number, private key, etc. A user may include user name, an email, a password, a role etc. An application may include application specifications. For example, a containerized system application (e.g., Kubernetes) may include YAML specifications. A data source may include abstraction for IoT sensors and gateways. A data pipeline may include metadata for a data pipeline, e.g., from input (e.g., data source or another data pipeline) through functions to output to cloud or edge data service. A function may include script or transformation, (e.g., Lambda runtime). A cloud profile may include a customer's cloud credential for cloud provider. A container registry profile may include a customer's container registry credential. A category may include name and values, used as a selector, e.g., Kubernetes labels for Kubernetes application.

(47) Generally, the CVMs 222(1)-(N) may be configured to control and manage any type of storage device of the storage 270. The CVMs 222(1)-(N) may implement storage controller logic and may virtualize all storage hardware of the storage 270 as one global resource pool to provide reliability, availability, and performance. IP-based requests may be generally used (e.g., by the user VMs 230(1)-(N) and/or the containers 232(1)-(N)) to send I/O requests to the CVMs 222(1)-(N). For example, the user VMs 230(1) and/or the containers 232(1) may send storage requests to the CVM 222(1) using an IP request, the user VMs 230(2) and/or the containers 232(2) may send storage requests to the CVM 222(2) using an IP request, etc. The CVMs 222(1)-(N) may directly implement storage and I/O optimizations within the direct data access path.

(48) Note that the CVMs 222(1)-(N) provided as virtual machines utilizing the hypervisors 210(1)-(N). Since the CVMs 222(1)-(N) run “above” the hypervisors 210(1)-(N), some of the examples described herein may be implemented within any virtual machine architecture, since the CVMs 222(1)-(N) may be used in conjunction with generally any type of hypervisor from any virtualization vendor.

(49) Virtual disks (vDisks) may be structured from the storage devices in the storage 270. A vDisk generally refers to the storage abstraction that may be exposed by the CVMs 222(1)-(N) to be used by the user VMs 230(1)-(N) and/or the containers 232(1)-(N). Generally, the distributed computing system 200 may utilize an IP-based protocol, such as an Internet small computer system interface (iSCSI) or a network file system interface (NFS), to communicate between the user VMs 230(1)-(N), the containers 232(1)-(N), the CVMs 222(1)-(N), and/or the hypervisors 210(1)-(N). Thus, in some examples, the vDisk may be exposed via an iSCSI or a NFS interface, and may be mounted as a virtual disk on the user VMs 230(1)-(N) and/or operating systems supporting the containers 232(1)-(N). iSCSI may generally refer to an IP-based storage networking standard for linking data storage facilities together. By carrying SCSI commands over IP networks, iSCSI can be used to facilitate data transfers over intranets and to manage storage over any suitable type of network or the Internet. The iSCSI protocol may allow iSCSI initiators to send SCSI commands to iSCSI targets at remote locations over a network. NFS may refer to an IP-based file access standard in which NFS clients send file-based requests to NFS servers via a proxy folder (directory) called “mount point”.

(50) During operation, the user VMs 230(1)-(N) and/or operating systems supporting the containers 232(1)-(N) may provide storage input/output (I/O) requests to the CVMs 222(1)-(N) and/or the hypervisors 210(1)-(N) via iSCSI and/or NFS requests. Each of the storage I/O requests may designate an IP address for a CVM of the CVMs 222(1)-(N) from which the respective user VM desires I/O services. The storage I/O requests may be provided from the user VMs 230(1)-(N) to a virtual switch within a hypervisor of the hypervisors 210(1)-(N) to be routed to the correct destination. For examples, the user VM 230(1) may provide a storage request to the hypervisor 210(1). The storage I/O request may request I/O services from a CVM of the CVMs 222(1)-(N). If the storage I/O request is intended to be handled by a respective CVM of the CVMs 222(1)-(N) hosted on a same respective computing node of the computing nodes 204(1)-(N) as the requesting user VM (e.g., CVM 222(1) and the user VM 230(1) are hosted on the same computing node

204(1)), then the storage I/O request may be internally routed within the respective computing node of the computing node of the computing nodes **204(1)-(N)**. In some examples, the storage I/O request may be directed to respective CVM of the CVMs **222(1)-(N)** on another computing node of the computing nodes **204(1)-(N)** as the requesting user VM (e.g., CVM **222(1)** is hosted on the computing node **204(1)** and the user VM **230(2)** is hosted on the computing node **204(2)**). Accordingly, a respective hypervisor of the hypervisors **210(1)-(N)** may provide the storage request to a physical switch to be sent over the network **280** to another computing node of the computing nodes **204(1)-(N)** hosting the requested CVM of the CVMs **222(1)-(N)**.

(51) The CVMs **222(1)-(N)** may collectively manage the storage I/O requests between the user VMs **230(1)-(N)** and/or the containers **232(1)-(N)** of the distributed computing system and a storage pool that includes the storage **270**. That is, the CVMs **222(1)-(N)** may virtualize I/O access to hardware resources within the storage pool. In this manner, a separate and dedicated CVM of the CVMs **222(1)-(N)** may be provided each of the computing nodes **204(1)-(N)** the distributed computing system **200**. When a new computing node is added to the distributed computing system **200**, it may include a respective CVM to share in the overall workload of the distributed computing system **200** to handle storage tasks. Therefore, examples described herein may be advantageously scalable, and may provide advantages over approaches that have a limited number of controllers. Consequently, examples described herein may provide a massively-parallel storage architecture that scales as and when computing nodes are added to the system.

(52) The distributed system **200** may include a centralized IoT manager that includes one or more of the centralized IoT manager instances **224(1)-(N)** hosted on the CVMs **222(1)-(N)**. The centralized IoT manager may be configured to centrally manage configuration of edge systems and data sources of the corresponding IoT system. In some examples, the centralized IoT manager may be configured to manage, for each of the edge systems, data sources, and/or users, network configuration and security protocols, installed software (e.g., including data pipelines and applications), connected data source(s) (e.g., including type, category, identifiers, data communication protocols, etc.), connected data plane(s), etc. The centralized IoT manager may maintain configuration information for each of the edge systems, data sources, associated users, including hardware configuration information, installed software version information, connected data source information (e.g., including type, category, identifier, etc.), associated data planes, current operational status, authentication credentials and/or keys, etc.

(53) In some examples, a workload of the centralized IoT manager may be distributed across two or more of the computing nodes **204(1)-(N)** via the respective centralized IoT manager instances **224(1)-(N)**. In other examples, the workload of the centralized IoT manager may reside in a single one of the centralized IoT manager instances **224(1)-(N)**. A number of centralized IoT manager instances **224(1)-(N)** running on the distributed computing system **200** may depend on a size of the management workload associated with the IoT system (e.g., based on a number of edge systems, data sources, users, etc., level of activity within the IoT system, frequency of updates, etc.), as well as compute resources available on each of the computing nodes **204(1)-(N)**. One of the centralized IoT manager instances **224(1)-(N)** may be designated a master centralized server manager that is configured to monitor workload of the centralized IoT manager instances **224(1)-(N)**, and based on the monitored workload, allocate management of respective edge systems and users to each of the centralized IoT manager instances **224(1)-(N)** and start additional centralized IoT manager instances when compute resources available to the centralized IoT manager have fallen below a defined threshold. Thus, while FIG. 2 depicts each of the CVMs **222(1)-(N)** hosting a respective one of the centralized IoT manager instances **224(1)-(N)**, it is appreciated that some of the CVMs **222(1)-(N)** may not have an active centralized IoT manager instances **224(1)-(N)** running without departing from the scope of the disclosure.

(54) In some examples, the centralized IoT manager may be configured to generate or update and distribute data pipelines and applications to selected edge systems based on the configuration maintained for each edge system. The data pipelines can be constructed using computing primitives and building blocks, such as VMs, containers, processes, or any combination thereof. In some examples, the data pipelines may be constructed using a group of containers (e.g., a pod) that each perform various functions within the data pipeline (e.g., subscriber, data processor, publisher, connectors that transform data for consumption by another container within the application or pod, etc.) In some examples, the definition of stages of a constructed data pipeline application may be described using a user interface or REST API, with data ingestion and movement handled by connector components built into the data pipeline. Thus, data may be passed between containers of a data pipeline using API calls.

(55) In some examples, an instance of the IoT manager instances **224(1)-(N)** may be respectively connected to a corresponding edge system of a plurality of edge systems. The IoT manager instance may further include a containerized system configured to receive, from a first instance of the IoT manager instances **224(1)-(N)**, a request to publish a message on a communication channel, and in response to the request to publish the message, send the message to the communication channel. The communication channel is established between a second instance of the IoT manager instances **224(1)-(N)** and a second edge system of the plurality of edge systems. Publishing of the message on the communication channel may cause the second instance of the IoT manager instances **224(1)-(N)** to send the message to the second edge system.

(56) The containerized system may be further configured to, upon establishment of the communication channel between the second instance of the IoT manager instances **224(1)-(N)** and the second edge system, receive a channel number of the communication channel from the second instance, and receive a look up request from the first instance of the IoT manager instances **224(1)-(N)** for the communication channel number associated with the second edge system. In response to the look up request, provide the channel number of the second edge system to the first instance of the IoT manager instances **224(1)-(N)**. The look up request may include an identity of the second edge system associated with the channel number. In some examples, the containerized system includes a publish/subscribe-to-channel map configured to store one or more communication channel numbers subscribed for receiving update notifications. The containerized system may also include an edge-to-channel map configured to store information associating edge systems to communication channels. In providing the channel number of the second edge system to the first instance of the IoT manager instances **224(1)-(N)**, the containerized system may use the edge-to-channel map to determine the channel number based on the edge system.

(57) In some examples, an instance of the IoT manager instances **224(1)-(N)** may be configured to: form secure two-way communication with an edge system of the IoT system, synchronize an object model with an edge object model for the edge system, and/or maintain the edge system using delta change communications. For example, an instance of the IoT manager instances **224(1)-(N)** may be configured to form a communication with the edge system of the IoT system, and detect a change of entity on the edge system based on an operation of the edge system. The instance of the IoT manager instances **224(1)-(N)**, may determine one or more additional edge systems that have associations with the detected change of entity, and transmit a respective message to the one or more additional edge systems. Each respective message may include information about a change to a corresponding edge system of the one or more additional edge systems. The instance of the IoT manager instances **224(1)-(N)** may also be configured to, in response to a request for inventory from the edge system, provide the inventory of the IoT system to the edge system. An edge system connected to an instance of

the IoT manager instances **224(1)-(N)** may receive the respective message from the instances of the IoT manager instances **224(1)-(N)**, update an inventory of the respective edge system based on the received respective message, and synchronize an application or a data pipeline of the respective edge system with the updated inventory.

(58) In some examples, a first instance of the IoT manager instances **224(1)-(N)** may be configured to receive a user message to deploy a data pipeline to an edge system connected to a second instance of the IoT manager instances **224(1)-(N)**, and determine a communication channel for the edge system. The first instance of the IoT manager instances **224(1)-(N)** may publish a message on the determined communication channel for the edge system. A second instance of the IoT manager instances **224(1)-(N)** may be configured to establish a communication channel with the edge system, and store a channel identification of the communication channel with association with the edge system at the IoT manager instances **224(1)-(N)** (e.g., the centralized IoT manager), and subscribe the communication channel to the IoT system data **272** for receiving update notifications. A second instance of the IoT manager instances **224(1)-(N)** may further receive the message published by the first instance of the IoT manager instances **224(1)-(N)**, and in response to receiving the message, send the message to the edge system.

(59) In some examples, an instance of the IoT manager instances **224(1)-(N)** may receive a request from an edge system for first authentication information, and provide the first authentication information to the new edge system. The instance of the IoT manager instances **224(1)-(N)** may receive a login request from the new edge system that includes second authentication information. The instance of the IoT manager instances **224(1)-(N)** may verify the second authentication information, and upon verification of the second authentication information, establish a secure communication channel with the new edge system.

(60) FIG. 3 is a block diagram of an example IoT system **300**, in accordance with an embodiment of the present disclosure. Examples described herein include a scalable centralized Internet of Things (IoT) manager of an IoT system capable of expanding as the number of tenants and edge systems within an IoT system increases. In some examples, an IoT system **300** may include tenants **302(1)-(N)**. Each of the tenants **302(1)-(N)** may include one or more respective edge systems. In other words, an edge system in the IoT system **300** may belong to a respective tenant of the tenants **302(1)-(N)**. Each of the tenants **302(1)-(N)** may have any number of edge systems, or have no edge systems. For example, a tenant of the tenants **302(1)-(N)** may include an edge cluster (e.g., **110** in FIG. 1), and an edge system may include any of the one or more of the edge cluster(s) **110**, the edge device(s) **112**, and/or the edge VM(s) **115** hosted on the server/cluster **114** of FIG. 1. In some examples, an edge system may include an edge device located in a tenant's private data center. Each edge device or system may include a containerized system (e.g., Kubernetes) cluster in itself. The IoT system **300** may also include a centralized IoT manager **342**. In some examples, each of the edge systems may be in communication with the centralized IoT manager **342** and federated by the centralized IoT manager **342**. The centralized IoT manager **342** may be operated in a cloud computing environment and provide a control plane for the federation of edge systems across multiple tenants.

(61) In some examples, an edge system of the tenants **302(1)-(N)** may include an edge stack. The edge stack may provide one or more services. In a non-limiting example, an edge stack may be configured to run one or more applications. For example, an edge stack may include a container as a service to run containerized system applications. By supporting known/off-the-shelf containerized system specifications (e.g., Kubernetes), the learning curve required of a user who wants to deploy an application may be reduced. The IoT system **300** may allow the user to define applications, such as via Kubernetes YAML specifications and control the deployment scope and lifecycle of such applications.

(62) In some examples, the edge stack may also include a data plane, wherein the data plane connects the edge system to the cloud. The data plane may also connect one or more sensors to an edge system. In a non-limiting example, the IoT system **300** may include a streaming data bus to facilitate data transfers between edge systems and the centralized IoT manager **342**. The IoT system **300** may also include built-in functions to allow a user to define custom functions and runtimes to transform the data along the data bus.

(63) The edge stack may store data in edge data services or send data to cloud data services of customer's cloud account. In some examples, the data may be sent from an edge to a cloud independent of the centralized IoT manager **342**. The communication between an edge system and the centralized IoT manager **342** is further explained.

(64) In some examples, the IoT manager may be configured to manage metadata of the system for each tenant and persist such metadata. The centralized IoT manager **342** may be an example implemented of the centralized IoT manager (e.g., the IoT manager **142** of FIG. 1, the centralized IoT manager formed by the IoT manager instances **224(1)-(N)** of FIG. 2). In some examples, the metadata of the IoT system **300** may be stored in IoT system data (e.g., the IoT system data **272** of FIG. 2). The metadata may be represented by object model, which classifies metadata into different entity. Examples of entities may include a tenant, an edge, a user, an application, a data source, a data pipeline, a function, a cloud profile, a container registry profile and/or a category. A tenant may include an instance per customer and tenant data key. A tenant may include one or more edge systems. An edge may include metadata for an edge system, such as serial number, private key, etc. A user may include user name, an email, a password, a role etc. An application may include application specifications. For example, a containerized system application (e.g., Kubernetes) may include YAML specifications. A data source may include abstraction for IoT sensors and gateways. A data pipeline may include metadata for a data pipeline, e.g., from input (e.g., data source or another data pipeline) through functions to output to cloud or edge data service. A function may include script or transformation, e.g., Lambda runtime. A cloud profile may include a customer's cloud credential for cloud provider. A container registry profile may include a customer's container registry credential. A category may include name and values, used as a selector, e.g., Kubernetes labels for Kubernetes application.

(65) In some examples, the centralized IoT manager **342** may include a set of APIs to expose its functionality. In some examples, the API may be a REST API. For example, the API may include a set of functions for each entity. In some examples, the set of functions for the API of the IoT manager may include a GET, PUT, POST and DELETE for each type of entity. For example, the centralized IoT manager **342** API may include POST/edges for creating an edge system having a payload that includes a representation of edge object, PUT/edges/<edge id> for updating an edge system having a payload that is an edge object. DELETE/edges/<edge id> for deleting an edge system, and GET/edges to retrieve all edge systems for the current tenant (e.g., a response is an array of edges). The various representations, such as edge object or array of edges, may be in any suitable format, such as JavaScript Object Notation (JSON). In some or other examples, the API of the centralized IoT manager **342** may include create, read, update and delete (CRUD) operations on the metadata entities.

(66) In some examples, the centralized IoT manager **342** may include a web-based user interface to enable a user to manage a respective IoT system **300**. For example, a user interface (UI), e.g., a web-based UI, may use REST API to communicate with the centralized IoT manager **342**. In some examples, the API (e.g., REST API) of the centralized IoT manager **342** may be available directly from the UI. In some examples, the REST API may also be accessible to one or more edge systems to synchronize with the centralized IoT manager

342.

(67) In some examples, the IoT system **300** may include communication links between each of the edge systems and the centralized IoT manager **342**. For example, the communication link may be a two-way channel (e.g., in WebSocket), which does not require any network setup other than the HTTPS outbound connection. In some examples, the centralized IoT manager **342** may be configured to monitor operations on an edge system and notify other edge systems of entity changes as a result of an operation on the edge system. For example, the centralized IoT manager **342** may monitor write operations on each of the edge system. When a write operation has occurred on an edge system, entity changes may occur on other edge systems. As such, the centralized IoT manager **342** may transmit a notification message, e.g., via a PUSH function, to other affected edge systems.

(68) In a non-limiting example, the communication link between an edge system and the centralized IoT manager **342** may include a messaging protocol (e.g., over a web socket notification). When an application is created or updated on an edge system, the centralized IoT manager **342** may transmit an onCreateApplication or an onUpdateApplication message to affected edge systems, where the affected edge systems may occur entity changes as a result of the write operation on the edge system. The payload for messages sent to the edge systems is the application created or updated. Similarly, when an application is deleted at an edge system, the centralized IoT manager **342** may send an onDeleteApplication message to affected edge systems, where the payload of the message is the Application ID.

(69) FIG. 4A is a flow diagram of an example process that may be implemented in an edge system to provide an update in an edge system in accordance with an embodiment of the present disclosure. In some examples, the process in FIG. 4A may be implemented in any of the one or more of the edge cluster(s) **110**, the edge device(s) **112**, and/or the edge VM(s) **115** hosted on the server/cluster **114** of FIG. 1 and/or any edge system of any of the tenants **302(1)-(N)** of FIG. 3. A process **400** may include booting up an edge system at **402**. Once an edge system boots up, if it is not yet connected to the IoT manager, the process **400** may establish a secure communication between the edge system and the IoT manager at operation **406** using an authentication method. The IoT manager may include the centralized IoT manager **142** of FIG. 1, the centralized IoT manager formed by the IoT manager instances **224(1)-(N)** of FIG. 2, the centralized IoT manager **342** of FIG. 3, or any combination thereof.

(70) FIG. 4B is a flow diagram of an example process that may be implemented in an edge device to establish a secure communication between an edge system and an IoT manager in accordance with an embodiment of the present disclosure. In some examples, a process **430** in FIG. 4B may be implemented in any of the one or more of the edge cluster(s) **110**, the edge device(s) **112**, and/or the edge VM(s) **115** hosted on the server/cluster **114** of FIG. 1 and/or any edge system of any of the tenants **302(1)-(N)** of FIG. 3. The IoT manager may include the centralized IoT manager **142** of FIG. 1, the centralized IoT manager formed by the IoT manager instances **224(1)-(N)** of FIG. 2, the centralized IoT manager **342** of FIG. 3, or any combination thereof.

(71) The process **430** may include receiving by an IoT manager, from a new edge system, a request for a first authentication information at operation **432**. In some examples, the new edge system may ping the IoT manager periodically using its serial number. In response, the process **430** may include, at the IoT manager, receiving a device serial number from an edge device, looking up the serial number of the edge device, obtaining the authentication information for the edge device based on the serial number, and providing the authentication information to the new edge system at operation **434**. For example, in response to the edge system requesting for authentication information, the IoT manager may return edge object metadata. The edge object metadata may include authentication information. The process **430** may further include, at the edge system, obtaining a private key at operation **436**. Once the private key is obtained, the process **430** may include, by the edge system, obtaining authentication information at operation **438**. For example, the edge system may invoke the/login REST API to obtain a secure token from the IoT manager. In response, the IoT manager may obtain the secure token and transmit the secure token to the edge device. In some examples, the token may be a JSON Web Token (JWT) that includes a set of information between the edge system and the IoT manager. For example, a JWT Token may include authentication information. Once the secure token is obtained, the process **430** may further include, by edge system, establishing a communication with the IoT manager using the secure token at operation **440**. In response, at operation **442**, the IoT manager may verify the secure token provided by the edge system, and if the secure token is verified, establish a safe communication channel with the new edge system.

(72) Returning to FIG. 4A, once the communication channel is established between the edge system and the IoT manager, the process **400** may include invoking APIs of the IoT manager to obtain inventory at **408**. The inventory may include connected data sources, data pipelines and applications, hardware resources, etc. or any combination thereof. Subsequently, the process **400** may create data pipeline at **410** and executing application at **412** based on the information contained in the inventory. Alternatively, and/or additionally, once an edge system is onboard, upon booting up, the process **400** may include synchronizing with the IoT manager to obtain the latest inventory at **414**. The process **400** may further sync up its application/data pipeline with the inventory at **416**.

(73) In some examples, when an edge system vice creates a new application/data pipeline or update an existing application or data pipeline, the IoT manager may determine all other edge systems within the tenant that need to receive change information. For example, an application running on one edge system may have access to one or more other edge systems within the tenant. When the application is deleted from the edge system, the associations of the application at each of the other edge systems will be affected. The IoT manager may notify these edge systems of the change of entity (e.g., application on the edge system). For example, the IoT manager may send a message to other affected edge systems within the tenant. In a non-limiting example, the message may be transmitted from the IoT manager to each of the affected edge systems via a connected web socket channel. In some examples, the payload for the message may include information about the change to the edge system (e.g., the deletion of an application).

(74) With further reference to FIG. 4A, the process **400** may include receiving a change message at **418**. The change message may be transmitted from the IoT manager in response to a change of entity on another edge system. The process **400** may further include updating the inventory at **420** based on the received change message. The process **400** may further include synchronizing the application/data pipeline state on the edge system with the updated inventory at **422**. Such operations ensure the edge system has the latest inventory.

(75) FIG. 5 is a block diagram of an IoT manager **542** of an IoT system **500** in accordance with an embodiment of the present disclosure. The IoT system may include the IoT manager **542** coupled to edge system(s) and user(s) **506(1)-(3)**, relational data base service (RDS) **582**, storage service **584**, and a search engine **586**. In some examples, the IoT manager **542** may be implemented in the centralized IoT manager **142** of FIG. 1, the centralized IoT manager formed by the IoT manager instances **224(1)-(N)** of FIG. 2, the centralized IoT manager **342** of FIG. 3, or any combination thereof. The edge system(s) of the edge system(s) and user(s) **506(1)-(3)** may each include any of the one or more of the edge cluster(s) **110**, the edge device(s) **112**, and/or the edge VM(s) **115** hosted on the server/cluster **114** of FIG. 1 and/or any edge system of any of the tenants **302(1)-(N)** of FIG. 3.

(76) In some examples, the IoT manager **542** may include one or more services. Examples of services include cloud manager service **590**, account service **591**, event service **592**, operator service **596**. In some examples, the cloud manager service **590** may be configured to store data, e.g., the APIs (e.g., REST API) for IoT manager **542**, user interface associated with API, configuration data for communication links (e.g., WebSocket) between IoT manager **542** and each of the edge systems. In some examples, the cloud manager service **590** may be configured to expose external endpoints to provide access to the external devices and services. The data stored in cloud manager may also include log bundles collected from edges. In some examples, the data in the cloud manager may be stored in the RDS **582**. For example, entities in the IoT system **500** may be represented by object models, and the object models for the entities are stored in the RDS **582**. In some examples, some data, e.g., log bundles collected from edges are stored in the storage service **584**, e.g., Amazon's simple storage service (S3). The storage service **584** may allow any suitable data to be stored and accessed via web interfaces, which allows scalability. In some examples, the cloud manager service **590** may have multiple instances to allow scalability, where an edge system may be connected to one of the instances, e.g., via WebSocket.

(77) In some examples, the account service **591** may be configured to store account related APIs in a similar manner as the cloud manager service **590**. For example, the data in the account service **591** may also be stored in the RDS **582** as object models. In some examples, the event service **592** may be configured to store data, e.g., event logging, application status, performance metrics and alerts. In a non-limiting example, data in the event service **592** may be stored and retrieved using any suitable database. Further, the event service **592** may use the search engine **586** to store and manage data. For example, the search engine may be implemented in Amazon's Elastic Search service. In a non-limiting example, the search engine **586** may store all logs from various the services **590-596** in the IoT manager **542**. The search engine **586** may also be queried by other edge systems and users of the edge system(s) and user(s) **506(1)-(3)**, e.g., via web interfaces. In some examples, the operator service **596** may be configured to store software update packages and provide edge software upgrade. As similar to cloud manager service **590**, the operator service **596** use storage service **584** (e.g., Amazon's S3) to store software update packages. The operator service **596** may also be configured to expose external endpoints to allow access by external devices and services.

(78) In some examples, the IoT manager **542** may include other services, such as secure service **594**, in-memory data service **593**, and/or tunnel service **595**. The secure service **594** may be configured to provide certificate management, such as SSL (secure socket layer) certificate management. In some examples, the secure service **594** may also use the RDS **582** to store tenant root certificate keys. In a non-limiting example, the secure service **594** may be implemented using CloudFlare SSL (CFSSL). The in-memory data service **593** may be configured to provide a key-value database to store suitable abstract data structure, such as text, bitmap, streams and indexes. The data service **593** may be a cache layer to store communication link (e.g., WebSocket) states and distributed caching. In a non-limiting example, the in-memory data service **593** may be implemented in various platforms, such as open source REDIS. The tunnel service **595** may be configured to provide tunnel that may connect servers behind a firewall to clients on the Internet, and this can be used to provide tunneling to an edge device. In some examples, the tunnel may be a HTTPS tunnel. The tunnel service **595** may be implemented in WStunnel available under Apache.

(79) In some examples, the above described services **590-596** in the IoT manager **542**, may include a containerized system service implemented by a containerized system **504** (e.g., Kubernetes) deployment. The containerized system **504** may also be configured to provide DNS service lookup, easy scale, incremental rollout, pod monitoring and auto restart, etc. One or more services **590-596** may communicate each other via a remote procedure call (e.g., RPC, gRPC).

(80) With further reference to FIG. 5, the IoT system **500** may also be configured to be scalable by using tenant-based sharding across multiple clusters, when the containerized system has a limit. In some examples, data services may be scaled up by using the RDS, ElasticSearch and/or S3 data services (from Amazon Web Services). In some examples, the S3 service may be used on demand, thus the load is not high. The IoT system **500** may use sharding on RDS or Elastic Search services. For example, the indices for Elastic Search may be sharded. Similarly, for RDS, the RDS instances may be scaled up by using tenant-based sharding.

(81) In some examples, the IoT system **500** may be configured to run multiple containerized system clusters (e.g., Kubernetes). For example, the IoT system **500** may include a developer cluster and a production cluster. The clusters may include multiple node cluster (e.g., 5 nodes) setup and use the containerized system's namespace to facilitate sharing of the cluster among developers. For example, the IoT system **500** may assign developers direct access through the APIs (e.g., kubectl access for Kubernetes). The IoT system **500** may provide each developer a respective namespace. The IoT system **500** may also provide each developer an instance of RDS DB and a DNS entry. The system may also provide each developer a UI universal resource locator (URL) (e.g., <https://<dev name>.containerizedsystem.com>). In some examples, the system may also include shared namespaces for testing. For example, the test namespace may run multiple cloud manager instances (e.g., **590**). The test name space may also run multiple account service instances (e.g., **591**).

(82) In some examples, the IoT system **500** may have multiple RDS instances (e.g., **582**). For example, the system may have one Writer instance and multiple Reader instances. The system may configure the Write instance to receive database queries, whereas the Reader instances may be configured to receive read calls. This reduces the load on the database Writer. In some examples, there may be a delay between the time Writer instance performs a writing operation and the time the Reader instance notes the latest write change. When the delay is significant, the Reader instance may read incorrect data.

(83) In some examples, the IoT system **500** may include a short lived key for write operations, where the short lived key has an expiration duration, e.g., 1 second, 2 seconds, 5 seconds etc. For example, when a write operation is activated, the system may set a short lived key for the expiration duration of time. After the expiration duration of time has passed, the system may automatically unset the short lived key. In a read operation, the system may check whether the short lived key is set or unset before invoking a database Reader to respond to a read query. If the short lived key is unset, the system may invoke the database Reader. If the short lived key is set, the system may wait before invoking the database Reader. The short lived key may be identified by <tenant id>/<entity type>. For example, a Read operation for a tenant/entity type may determine the absence of <tenant id>/<entity type> key (when such absence indicates no recent writes for the tenant/entity) before using DB reader to serve the query.

(84) Alternatively, and/or additionally, the IoT system **500** may determine whether database inconsistency exists. For example, the system may determine that the delay between write and read of the latest write becomes significant. If the system determines that the database inconsistency exists, the system may decide not to create multiple Reader instances.

(85) In some examples, one or more services or systems disclosed in the present disclosure (e.g., **590-596**) may be configured to store user login failure count, e.g., in a REDIS key. For example, in each failed login for email, an email service may increment the error count. When the error count has exceed a threshold (e.g., 5), the email service may block user login. Additionally, and/or alternatively,

the email service may also limit the number of login attempts to a threshold count per unit time period, e.g., 5 times per minute. In such a way, the email service may effectively make brute force attack to guess user login password more difficult.

(86) In some examples, the IoT system **500** may implement a role-based access control (RBAC) to provide protections for the system to be scalable to manage larger organizations. In some examples, full-fledged RBAC with fine grained user roles and permissions may be implemented. In other examples, the system may implement project-based RBAC, in which the system may facilitate all users to gain access to data entities in a project and infra entities associated with a project via project assignment. For example, the IoT system **500** may assign a user zero or more projects. The system may also create data entities, such as applications, data pipelines, functions, etc. within a project in response to a user request, where the user must also have access to the project. Entities, such as edge devices, cloud profiles, container registry profiles, etc., that are associated with the applications, data pipelines, functions etc., may also be associated with zero or more projects via project creation, update, or deletion. The system may further classify users into admin users and non-admin users, in which only admin users have write access to projects.

(87) In implementing project-based RBAC, the object model disclosed in the present disclosure for each tenant may be updated with projects. In some examples, data entities will each have a project_id attribute indicating which project it belongs to. Infra entities associated with data entities may include an entity-project association table to capture the many-to-many relationship between the entities and projects. In some examples, for each entity type, its CRUD API behavior may be specified under RBAC. For example, for an application entity, CUD API is allowed if and only if the user is assigned to the project associated with the entity. A Read API for an application entity will return all applications in all projects the user is assigned to. In another example, for a cloud profile entity, CUD API is allowed if and only if the user is an admin user. A call to a Read API for cloud profile from a non-admin user will return all cloud profiles associated with all projects the user is assigned to. A call to a Read API for cloud profile from an admin user, on the other hand, will return all cloud profiles regardless of project association. This is because admin users need to manage all cloud profiles. Although the examples only show application entity and cloud profile entity, other entity type, operation type, user type, user project membership, or combinations thereof may also be implemented.

(88) The various embodiments disclosed in the present disclosure may allow various communications in the IoT system. FIG. 6 is a block diagram illustrating communication via a centralized IoT manager **642** of an IoT system **600**, in accordance with an embodiment of the present disclosure. The IoT manager **642** may serve all edge systems (e.g., any of the one or more of the edge cluster(s) **110**, the edge device(s) **112**, and/or the edge VM(s) **115** hosted on the server/cluster **114** of FIG. 1, any edge system of any of the tenants **302(1)-(N)** of FIG. 3, any of the edge system(s) of the edge system(s) and user(s) **506(1)-(3)** of FIG. 5, or combinations thereof) and users (e.g., any of the user(s) of the edge system(s) and user(s) **506(1)-(3)** of FIG. 5) across all tenants of the IoT system **600**. The IoT manager **642** may be implemented in the centralized IoT manager **142** of FIG. 1, the centralized IoT manager formed by the IoT manager instances **224(1)-(N)** of FIG. 2, the centralized IoT manager **342** of FIG. 3, IoT manager **542** of FIG. 5, or any combination thereof. Various configurations of the IoT manager **642** are included to meet scalability requirements. For example, the IoT manager **642** may support hundreds or thousands of customers, with over 100 edge systems each. Each edge system may make up to dozens of requests per minute. In some examples, the IoT manager **642** may have first and second cloud manager instances **690(1)-(2)** each connected to an edge system, **606(1)-606(2)**, respectively. Although FIG. 6 shows two instances of cloud manager services, other number of instances may be possible. While FIG. 6 depicts the two cloud manager instances **690(1)-(2)** each connected to a single one of the edge systems **606(1)-(2)**, each of the cloud manager instances **690(1)-(2)** may be connected to more than one respective edge system. The cloud manager instances **690(1)-(2)** may be in communication with the containerized system **693**. The containerized system **693** may be implemented in a REDIS (e.g., by Amazon). At least one of the cloud manager instances **690(1)-(2)** may be accessible by a user **605**, e.g., to receive and send a message from/to the user.

(89) In some examples, the IoT manager **642** may be scalable by externalizing states of services (e.g., **590-596** in FIG. 5) into data services so that these services may become stateless. For example, the IoT manager **642** may externalize the WebSocket connection states among the edge systems **606(1)-(2)** and the cloud manager instances **690(1)-(2)**. In some examples, the containerized system **693** may include an edge-to-channel mapping **698**. The edge-to-channel mapping **698** may be configured to cache the edge to cloud manager relationships. For example, the first edge system **606(1)** is connected to the first cloud manager instance **690(1)** using a first communication channel "ID 1", and the second edge system **606(2)** may be connected to the second cloud manager instance **690(2)** using a second communication channel "ID 2." The containerized system **693** is further configured to allow any of the cloud manager instances **690(1)-(2)** to look up the edge-to-channel mapping **698**. For example, one of the cloud manager instances **690(1)-(2)** may query the containerized system **693** for a channel number given the edge system ID.

(90) In some examples, the containerized system **693** may further include a publish/subscribe-to-channel mapping **697**. The publish/subscribe-to-channel mapping **697** may be configured to cache one or more communication channels that have subscribed to the publishing of messages related to the edge system(s) **606(1)-(2)** on each respective communication channel. An instance of the cloud manager instances **690(1)-(2)** on a communication channel may be configured to receive a publish message sent by another instances of the cloud manager instances **690(1)-(2)**. A publish message may include information about entities (e.g., edge systems **606(1)-(2)**) and change of entities. Upon receiving a message in a communication channel, an instance of the cloud manager instances **690(1)-(2)** in the communication channel may operate upon one or more edge systems indicated in the publish message.

(91) With further reference to FIG. 6, an example operation of the containerized system **693** and messaging between the edge systems **606(1)-(2)** and the cloud manager instances **690(1)-(2)** is illustrated. In some examples, a first edge system **606(1)** is connected to the first cloud manager instance **690(1)** in a first communication channel CH ID 1. A second edge system **606(2)** is connected to the second cloud manager instance **690(2)** in a second communication channel CH ID 2 at operation (1). The second cloud manager instance **690(2)** may subscribe to messages on the second communication channel CH ID 2 at operation (2). The containerized system **693** may cache the second edge system **606(2)** to second communication channel CH ID 2 relationship in the edge-to-channel mapping **698**. For example, the second cloud manager instance **690(2)** may add the second edge system **606(2)** in association with the second communication channel CH ID 2 in the edge-to-channel mapping **698** at operation (3). An edge system may be identified by an edge device ID.

(92) In some examples, the containerized system **693** may further store information about channel(s) subscribed in the publish/subscribe-to-channel mapping **697**. For example, upon the second cloud manager instance **690(2)** having subscribed to channel CH ID 2, the publish/subscribe-to-channel mapping **697** may store the subscribed channel number CH ID 2. The containerized system **693** may further be configured to receive a request from any of the cloud manager instances **690(1)-(2)** to publish a message on a given communication channel, and in response, relay the message on the provided communication channel.

(93) With further reference to FIG. 6, the user **605** may send a message to the first cloud manager instance **690(1)** at operation (4). The message may be sent to the cloud manager instance **690(1)** via an IoT manager API call, e.g., a REST API call. In some examples, the message may include a write operation, which may invoke other operations on one or more of the edge systems **606(1)-(2)**. For example, a user command in the user message may include deploying a data pipeline to one or more of the edge systems **606(1)-(2)**. In this example, the edge system **606(1)** is directly connected to the first cloud manager instance **690(1)** receiving the message from the user **605**, while the edge system **606(2)** is connected to the second cloud manager instance **690(2)**. In the instant example, the write operation may deploy data pipelines to both the first edge system **606(1)** and the second edge system (**606(2)**), which are connected to the first and second cloud manager instances **690(1)-(2)**, respectively.

(94) With continued reference to FIG. 6, upon receiving the message from the user **605**, the first cloud manager instance **609(1)** may send a data pipeline update message to the first edge system **606(1)** at operation (5), e.g., via a WebSocket message, because the first edge system **606(1)** is directly connected to the first cloud manager instance **609(1)**. The payload of the message may be the REST message sent by the user at operation (4). For example, the message may contain parameters associated with the data pipeline change. In response to receiving the message, the first edge system **606(1)** may implement the data pipeline change specified in the message.

(95) With further reference to FIG. 6, in response to receiving the user message at operation (4), the first cloud manager instance **609(1)** may further determine the communication channel associated with the second edge system **606(2)** at operation (6). For example, the first cloud manager instance **609(1)** may look up the edge to channel mapping **698** in the containerized system **693** by the device ID of the second edge system **606(2)**. In the instant example, the determined communication channel is CH ID 2.

(96) Upon determining the communication channel(s) associated with the edge system **606(2)** related to the message from the user **605**, the first cloud manager instance **690(1)** may further publish the message on the second communication channel CH ID 2 at operation (6). In response to the publishing of the message at the second communication channel, the second cloud manager instance **609(2)** receives the message at operation (7) and subsequently send the message to the second edge system **606(2)** on the second communication channel at operation (8). The message sent to the second edge system **606(2)** may contain parameters associated with data pipeline change regarding the second edge system **606(2)**. In response to receiving the message, the second edge system **606(2)** may perform the operations as specified in the message. When the containerized system **693** is implemented in REDIS, the cloud manager instances **690(1)-(2)** may publish the messages in a REDIS message.

(97) The above configuration of the IoT system **600** may allow scalability of the system. For example, the limit of the containerized system **693** may be the size of the memory. As such, the containerized system **693** may readily support **100** subscribers and **100,000** messages, and handle messages by over **10,000** customers. Although operations in the above example are illustrated in numerals, it is appreciated that the order of the operations may vary. Further, a subset of the illustrated operations may also be independently performed without requiring all of the operations to be performed.

(98) FIG. 7 is a flow diagram of a method **700** to communicate with an edge system of an IoT system, in accordance with an embodiment of the present disclosure. The method **700** may be implemented in an IoT system, such as the IoT system **100** of FIG. 1, the IoT system **200** of FIG. 2, the IoT system **300** of FIG. 3, the IoT system of **500** of FIG. 5, the IoT system **600** of FIG. 6, of a combination thereof. The method **700** may include forming a secure two-way communication with an edge system of an IoT system via a centralized IoT manager hosted on a computing system at operation **710**. The centralized IoT manager may be formed by the IoT manager instances **224(1)-(N)** of FIG. 2, the centralized IoT manager **342** of FIG. 3, or any combination thereof. The edge system may include any of the edge systems shown in FIGS. 1, 3, 5 and 6. In some examples, the centralized IoT manager may support two-way communication using secure protocols, such as HTTPS, WSS protocol, etc. A metadata object model may be provided to fully specify behavior of edge systems.

(99) In some examples, the method **700** may further synchronize an object model at the centralized IoT manager with an edge object model for the edge system at operation **720**. This is referred to as initial synchronization. In some examples, the object model of the edge system may persist at the edge system to allow the edge system to continue to operate offline. The method **700** may further maintain the edge system using a delta change communication and notification at operation **730**. For example, the IoT manager may manage subsequent object model synchronizations after the initial synchronization using delta-based communication and update notification (e.g., WebSocket), which is more efficient than full object model synchronizations. For example, the IoT manager may monitor the changes of the object model of the IoT system and/or the changes of the object model for the edge system, and synchronize the object model of the system and the object model of the edge system only with respect to the changes of each respective object model. The method **700** may be implemented in any suitable manner. For example, the method **700** may implement the process **400, 430** (in FIGS. 4A and 4B) or a subset of the processes **400** and **430**.

(100) FIG. 8 is a flow diagram of a method **800** to provide software update to an edge system of an IoT system, in accordance with an embodiment of the present disclosure. The method **800** may be implemented in an IoT system, such as the IoT system **100** of FIG. 1, the IoT system **200** of FIG. 2, the IoT system **300** of FIG. 3, the IoT system of **500** of FIG. 5, the IoT system **600** of FIG. 6, of a combination thereof. The method **800** may include receiving, at a centralized IoT manager of an IoT system, a new version of an application hosted on an edge system of the IoT system at operation **810**. In some examples, the centralized IoT manager of the IoT system may be configured to centrally manage edge systems of the IoT system, including pushing software and/or configuration updates to the edge systems, providing data pipelines and other applications to the edge systems, storing/updating respective metadata for the IoT system in a database, etc. The centralized IoT manager may be formed by the IoT manager instances **224(1)-(N)** of FIG. 2, the centralized IoT manager **342** of FIG. 3, or any combination thereof. The edge system may include any of the edge systems shown in FIGS. 1, 3, 5 and 6.

(101) The method **800** may further determine a delta change between the new version and the current version of the application at operation **820**. In some examples, the method **800** may compare the object model of the system with the object model of the edge system that hosts the application.

(102) With further reference to FIG. 8, the method **800** may further establish a secure communication channel with the edge system at operation **830**. Various communication protocols may be possible. For example, the method **800** may establish the communication channel using a WebSocket communication protocol. Upon establishing the communication channel between the IoT manager and the edge system, the method **800** may provide the delta change between the new version and the current version of the application to the edge system at operation **840**. In some examples, the centralized IoT manager may manage the software updates of the edge systems by pushing the software updates according to a determined schedule managed by the centralized IoT manager. In other words, one or more of operations **810, 820, 830** and **840** may be configured to perform at the determined schedule. For example, the method **800** may

periodically check for software update and receive the new version of the software at operation **810**. In some examples, the method **800** may be implemented in any suitable manner. For example, the method **800** may implement the process **400**, **430** (in FIGS. **4A** and **4B**) or a subset of the processes **400** and **430**.

(103) FIG. **9** depicts a block diagram of components of an edge system and/or a computing node (device) **900** in accordance with an embodiment of the present disclosure. It should be appreciated that FIG. **9** provides only an illustration of one implementation and does not imply any limitations with regard to the environments in which different embodiments may be implemented. Many modifications to the depicted environment may be made. The device **900** may implemented as any of an edge system of the edge cluster(s) **110**, the edge device(s) **112**, the server/cluster **114**, a computing node of the central IoT computing system **140**, or a computing node of the data computing system **150** of FIG. **1**, all or part of the edge computing system **200** of FIG. **2**, any of the computing nodes **204(1)-(N)** of FIG. **2**, any of the edge systems of the tenants **302(1)-(N)** and/or a computing node configured to host the centralized IoT manager **342** of FIG. **3**, any of the edge systems of the edge(s) and user(s) **506(1)-(3)** and/or any device configured to host the centralized IoT manager **542** of FIG. **5**, either of the edge systems **606(1)-(2)** and/or devices configured host the centralized IoT manager **642** of FIG. **6**, or any combination thereof. The device **900** may be configured to implement the method **700** of FIG. **7**, and method **800** of FIG. **8**, or any combination thereof.

(104) The device **900** includes a communications fabric **902**, which provides communications between one or more processor(s) **904**, memory **906**, local storage **908**, communications unit **910**, I/O interface(s) **912**. The communications fabric **902** can be implemented with any architecture designed for passing data and/or control information between processors (such as microprocessors, communications and network processors, etc.), system memory, peripheral devices, and any other hardware components within a system. For example, the communications fabric **902** can be implemented with one or more buses.

(105) The memory **906** and the local storage **908** are computer-readable storage media. In this embodiment, the memory **906** includes random access memory RAM **914** and cache **916**. In general, the memory **906** can include any suitable volatile or non-volatile computer-readable storage media. The local storage **908** may be implemented as described above with respect to local storage **206(1)-(N)** and/or storage **270** of FIG. **2**. In this embodiment, the local storage **908** includes an SSD **922** and an HDD **924**, which may be implemented as described above with respect to any of SSD **240(1)-(N)** and any of HDD **242(1)-(N)**, respectively.

(106) Various computer instructions, programs, files, images, etc. may be stored in local storage **908** for execution by one or more of the respective processor(s) **904** via one or more memories of memory **906**. In some examples, local storage **908** includes a magnetic HDD **924**. Alternatively, or in addition to a magnetic hard disk drive, local storage **908** can include the SSD **922**, a semiconductor storage device, a read-only memory (ROM), an erasable programmable read-only memory (EPROM), a flash memory, or any other computer-readable storage media that is capable of storing program instructions or digital information.

(107) The media used by local storage **908** may also be removable. For example, a removable hard drive may be used for local storage **908**. Other examples include optical and magnetic disks, thumb drives, and smart cards that are inserted into a drive for transfer onto another computer-readable storage medium that is also part of local storage **908**.

(108) Communications unit **910**, in these examples, provides for communications with other data processing systems or devices. In these examples, communications unit **910** includes one or more network interface cards. Communications unit **910** may provide communications through the use of either or both physical and wireless communications links.

(109) I/O interface(s) **912** allows for input and output of data with other devices that may be connected to device **900**. For example, I/O interface(s) **912** may provide a connection to external device(s) **918** such as a keyboard, a keypad, a touch screen, and/or some other suitable input device. External device(s) **918** can also include portable computer-readable storage media such as, for example, thumb drives, portable optical or magnetic disks, and memory cards. Software and data used to practice embodiments of the present disclosure can be stored on such portable computer-readable storage media and can be loaded onto local storage **908** via I/O interface(s) **912**. I/O interface(s) **912** also connect to a display **920**. Display **920** provides a mechanism to display data to a user and may be, for example, a computer monitor.

(110) Various features described herein may be implemented in hardware, software executed by a processor, firmware, or any combination thereof. If implemented in software (e.g., in the case of the methods described herein), the functions may be stored on or transmitted over as one or more instructions or code on a computer-readable medium. Computer-readable media includes both non-transitory computer storage media and communication media including any medium that facilitates transfer of a computer program from one place to another. A non-transitory storage medium may be any available medium that can be accessed by a general purpose or special purpose computer. By way of example, and not limitation, non-transitory computer-readable media can comprise RAM, ROM, electrically erasable programmable read only memory (EEPROM), or optical disk storage, magnetic disk storage or other magnetic storage devices, or any other non-transitory medium that can be used to carry or store desired program code means in the form of instructions or data structures and that can be accessed by a general-purpose or special-purpose computer, or a general-purpose or special-purpose processor.

(111) From the foregoing it will be appreciated that, although specific embodiments of the disclosure have been described herein for purposes of illustration, various modifications may be made without deviating from the spirit and scope of the disclosure. Accordingly, the disclosure is not limited except as by the appended claims.

Claims

1. At least one non-transitory computer-readable storage medium including instructions that, when executed, cause a system to: store a map associating multiple edge systems with respective communication channels, wherein each of the communication channels are associated with a respective instance of a centralized manager, the instances of the centralized manager forming a distributed centralized manager providing a control plane for the edge systems; access the map to identify a selected communication channel of the communication channels, the selected communication channel associated with a specified edge system, the specified edge system including an edge stack having a machine learning (ML) service configured to provide inference data; publish an update message using the control plane on the selected communication channel to update the specified edge system including the ML service; and provide the inference data to a data plane of a computing system.
2. The at least one non-transitory computer-readable storage medium of claim 1, the instructions, when executed, further cause the system to: store a map associating one or more channel numbers with one or more communication channels, including the respective

communication channels, the one or more channel numbers indicative of which of the one or more communication channels are subscribed to publishing of messages associated with the edge systems, including update messages.

3. The at least one non-transitory computer-readable storage medium of claim 2, wherein providing the inference data to the data plane of the computing system is based at least on using a data plane communication interface, including an application programming interface, wherein the data plane communication interface is configured to communicate with one or more other data planes.

4. The at least one non-transitory computer-readable storage medium of claim 1, the instructions, when executed, further cause the system to: scale the instances of the centralized manager to increase based at least on a determination that a number of tenants, a number of edge systems, or a combination thereof, increases.

5. The at least one non-transitory computer-readable storage medium of claim 1, wherein the respective instance of the centralized manager is communicatively coupled to the specified edge system including the edge stack having the ML service, via a WebSocket.

6. The at least one non-transitory computer-readable storage medium of claim 1, the instructions, when executed, further cause the system to: receive a request, at the ML inference service, for an ML model inference; based at least on receipt of the request, load the requested ML model inference in an inference engine; output the inference data; and provide the inference data to the data plane of the computing system.

7. The at least one non-transitory computer-readable storage medium of claim 1, wherein the data plane connects one or more sensors to the edge system.

8. The at least one non-transitory computer-readable storage medium of claim 1, wherein the data plane connects the respective edge system to a cloud.

9. The at least one non-transitory computer-readable storage medium of claim 1, wherein the data plane is configured to store the inference data, process the inference data, provide access to the inference data, or combinations thereof.

10. A method comprising: storing a map associating multiple edge systems with respective communication channels, wherein each of the communication channels are associated with a respective instance of a centralized manager, the instances of the centralized manager forming a distributed centralized manager providing a control plane for the edge systems; accessing the map to identify a selected communication channel of the communication channels, the selected communication channel associated with a specified edge system, the specified edge system including an edge stack having a machine learning (ML) service configured to provide inference data; publishing an update message using the control plane on the selected communication channel to update the specified edge system including the ML service; and providing the inference data to a data plane of a computing system.

11. The method of claim 10, further comprising storing a map associating one or more channel numbers with one or more communication channels, including the respective communication channels, the one or more channel numbers indicative of which of the one or more communication channels are subscribed to a publishing of messages associated with the edge systems, including update messages.

12. The method of claim 11, wherein providing the inference data to the data plane of the computing system is based at least on using a data plane communication interface, including an application programming interface, wherein the data plane communication interface is configured to communicate with one or more other data planes.

13. The method of claim 10, further comprising scaling the instances of the centralized manager to increase based at least on a determination that a number of tenants, a number of edge systems, or a combination thereof, increases.

14. The method of claim 10, wherein the respective instance of the centralized manager is communicatively coupled to the specified edge system including the edge stack having the ML service, via a WebSocket.

15. The method of claim 10, further comprising: receiving a request, at the ML inference service, for an ML model inference; based at least on receipt of the request, loading the requested ML model inference in an inference engine; outputting the inference data; and providing the inference data to the data plane of the computing system.

16. The method of claim 10, wherein the data plane connects one or more sensors to the edge system.

17. The method of claim 10, wherein the data plane connects the respective edge system to a cloud.

18. The method of claim 10, wherein the data plane is configured to store the inference data, process the inference data, provide access to the inference data, or combinations thereof.

19. A system comprising: a central computing system comprising at least one processor and memory storing instructions which when executed by the at least one processor, cause a respective instance of a centralized manager to perform operations, the central computing system hosting one or more instances of the centralized manager, including the respective instance of the centralized manager, the one or more instances of the centralized manager forming a distributed centralized manager providing a control plane for one or more edge systems; and the one or more edge systems, wherein the respective instance of the centralized manager is configured to: store a map associating the one or more edge systems with respective communication channels, wherein each of the communication channels are associated with respective instances of the centralized manager; access the map to identify a selected communication channel of the communication channels, the selected communication channel associated with a specified edge system of the one or more edge systems, wherein the specified edge system includes an edge stack having a machine learning (ML) service configured to provide inference data; publish an update message using the control plane on the selected communication channel to update the specified edge system including the ML service; and provide the inference data to a data plane of a computing system.

20. The system of claim 19, wherein the respective instance of the centralized manager is further configured to store a map associating one or more channel numbers with one or more communication channels, including the respective communication channels, the one or more channel numbers indicative of which of the one or more communication channels are subscribed to a publishing of messages associated with the one or more edge systems, including update messages.

21. The system of claim 20, wherein providing the inference data to the data plane of the computing system is based at least on using a data plane communication interface, including an application programming interface, wherein the data plane communication interface is configured to communicate with one or more other data planes.

22. The system of claim 19, wherein the central computing system is further configured to scale the instances of the centralized manager to increase based at least on a determination that a number of tenants, a number of edge systems, or a combination thereof, increases.

23. The system of claim 19, wherein the respective instance of the centralized manager communicatively coupled to the specified edge system including the edge stack having the ML service, via a WebSocket.

24. The system of claim 19, further comprising: provide a request, by the respective instance of the centralized manager and to the ML inference service, for an ML model inference, the request causing the ML inference service to: load the requested ML model inference in

an inference engine, output the inference data, and provide the inference data to the data plane of the computing system.

25. The system of claim 19, wherein the data plane connects one or more sensors to the edge system.

26. The system of claim 19, wherein the data plane connects the respective edge system to a cloud.

27. The system of claim 19, wherein the data plane is configured to store the inference data, process the inference data, provide access to the inference data, or combinations thereof.
