

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent

12399785

Kind Code

B2

Date of Patent

August 26, 2025

Inventor(s)

Mitkar; Amit et al.

Backup of containerized applications using a backup services container and a backup services container-orchestration pod

Abstract

A “backup services container” comprises “backup toolkits,” which include scripts for accessing containerized applications plus enabling utilities/environments for executing the scripts. The backup services container is added to Kubernetes pods comprising containerized applications without changing other pod containers. For maximum value and advantage, the backup services container is “over-equipped” with toolkits. The backup services container selects and applies a suitable backup toolkit to a containerized application to ready it for a pending backup. Interoperability with a proprietary data storage management system provides features that are not possible with third-party backup systems. Some embodiments include one or more components of the proprietary data storage management within the illustrative backup services container. Some embodiments include one or more components of the proprietary data storage management system in a backup services pod configured in a Kubernetes node. All configurations and embodiments are suitable for cloud and/or non-cloud computing environments.

Inventors: Mitkar; Amit (Manalapan, NJ), Degaonkar; Sumedh Pramod (Sammamish, WA), Kumar; Sanjay (Morganville, NJ), Bhavanarushi; Shankarbabu (Hyderabad, IN), Kumar; Vikash (Edison, NJ)

Applicant: Commvault Systems, Inc. (Tinton Falls, NJ)

Family ID: 1000008782106

Assignee: Commvault Systems, Inc. (Tinton Falls, NJ)

Appl. No.: 18/731501

Filed: June 03, 2024

Prior Publication Data

Document Identifier

Publication Date

US 20240330120 A1

Oct. 03, 2024

Related U.S. Application Data

continuation parent-doc US 18074787 20221205 US 12007854 child-doc US 18731501
continuation parent-doc US 16924002 20200708 US 11561866 20230124 child-doc US 18074787
us-provisional-application US 62872606 20190710

Publication Classification

Int. Cl.: G06F11/14 (20060101); G06F9/455 (20180101); G06F9/50 (20060101)

U.S. Cl.:

CPC G06F11/1464 (20130101); G06F9/45558 (20130101); G06F9/5077 (20130101);
G06F11/1451 (20130101); G06F11/1458 (20130101); G06F11/1461 (20130101);
G06F11/1469 (20130101); G06F11/1484 (20130101); G06F2009/45562 (20130101);
G06F2009/45583 (20130101); G06F2201/84 (20130101)

Field of Classification Search

CPC: G06F (9/4881); G06F (9/5038)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2008/0244177	12/2007	Crescenti	711/E12.001	G06F 11/1458

OTHER PUBLICATIONS

“Production-Grade Container Orchestration”, Kubernetes , accessed on <https://kubernetes.io/>, Jul. 10, 2019 available on <https://web.archive.org/web/20190710195441/http://kubernetes.io/>, May 3, 2024, 8 pages. cited by applicant

Google Cloud Regions and Zones, accessed on <https://cloud.google.com/compute/docs/regions-zones/>, Apr. 26, 2019, available on <https://web.archive.org/web/20190415102759/cloud.google.com/compute/docs/regions-zones/>, Dec. 4, 2023, 30 pages. cited by applicant

Kubernetes kubelet, Kubernetes Documentation, accessed on <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/>, May 2020, available on https://web.archive.org/web/20200510051539mp_/https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/, Mar. 29, 2024, 13 pages. cited by applicant

Margaret Rouse, “application containerization (app containerization)”, TechTarget, accessed on <https://www.techtarget.com/searchitoperations/definition/application-containerization-app-containerization>, Jul. 2019, available on <https://web.archive.org/web/20190703014111/http://searchitoperations.techtarget.com/definition/application-containerization-app-containerization>, Mar. 29, 2024, 7 pages. cited by applicant

Margaret Rouse, “Definition of Availability Zones”, TechTarget, accessed on searchaws.techtarget.com/definition/availability-zones, Apr. 26, 2019, available on <https://web.archive.org/web/20180911194556/https://searchaws.techtarget.com/definition/availability-zones>, Dec. 4, 2023, 5 pages. cited by applicant

Mell et al., The NIST Definition of Cloud Computing, 800-145, Sep. 2011, 7 pages. cited by applicant

Velero, “Backup and migrate Kubernetes applications and their persistent volumes”, accessed on

<https://github.com/vmware-tanzu/velero>, May 3, 2024, 4 pages. cited by applicant

Viewing Pods and Nodes, Kubernetes Documentation, accessed on <https://kubernetes.io/docs/tutorials/kubernetes-basics/explore/explore-intro/>, 2021, available on <https://web.archive.org/web/20210622203553/http://kubernetes.io/docs/tutorials/kubernetes-basics/explore/explore-intro/>, Mar. 29, 2024, 3 pages. cited by applicant

Watts et al. “SaaS vs PaaS vs IaaS: What's The Difference & How To Choose”, BMC Blogs, BMC Software, Inc., accessed on <https://www.bmc.com/blogs/saas-vs-paas-vs-iaas-whats-the-difference-and-how-to-choose/>, Jun. 15, 2019, 18 Pages. cited by applicant

What is Kubernetes?, Kubernetes Documentation, accessed on <https://kubernetes.io/docs/concepts/architecture/nodes/>, Jul. 2021, available on <https://web.archive.org/web/20210724121336/https://kubernetes.io/docs/concepts/overview/what-is-kubernetes/>, Mar. 29, 2024, 4 pages. cited by applicant

Wikipedia, “Docker (software)”, accessed on [https://en.wikipedia.org/wiki/Docker_\(software\)](https://en.wikipedia.org/wiki/Docker_(software)), 2019 , available on [https://web.archive.org/web/20190403180753/https://en.wikipedia.org/wiki/Docker_\(software\)](https://web.archive.org/web/20190403180753/https://en.wikipedia.org/wiki/Docker_(software)), May 27, 2024, 9 pages. cited by applicant

Wikipedia, “Kubernetes—Volumes”, accessed on <https://en.wikipedia.org/wiki/Kubernetes#Volumes>, Jul. 2021, available on <https://web.archive.org/web/20210723152834/http://en.wikipedia.org/wiki/Kubernetes#Volumes>, Mar. 29, 2024, 16 pages. cited by applicant

Wikipedia, “Kubernetes”, accessed on <https://en.wikipedia.org/wiki/Kubernetes>, May 3, 2024, 24 pages. cited by applicant

Primary Examiner: Rashid; Wissam

Attorney, Agent or Firm: Commvault Systems, Inc.

Background/Summary

INCORPORATION BY REFERENCE TO ANY PRIORITY APPLICATIONS (1) This application is a Continuation of U.S. patent application Ser. No. 18/074,787 filed on Dec. 5, 2022, which claims priority to U.S. patent application Ser. No. 16/924,002 filed on Jul. 8, 2020, which claims priority to U.S. Provisional Patent Application Ser. No. 62/872,606 filed on Jul. 10, 2019, which is incorporated by reference in its entirety herein. Any and all applications for which a foreign or domestic priority claim is identified in the Application Data Sheet of the present application are hereby incorporated by reference in their entireties under 37 CFR 1.57.

COPYRIGHT NOTICE

(1) A portion of the disclosure of this patent document contains material which is subject to copyright protection. The copyright owner has no objection to the facsimile reproduction by anyone of the patent document and/or the patent disclosure as it appears in the United States Patent and Trademark Office patent file and/or records, but otherwise reserves all copyrights whatsoever.

BACKGROUND

(2) Businesses recognize the commercial value of their data and seek reliable, cost-effective ways to protect the information stored on their computer networks while minimizing impact on productivity. A company might back up critical computing systems such as databases, file servers, web servers, virtual machines, and so on as part of a maintenance plan. Given the rapidly expanding volume of data under management, companies also continue to seek innovative techniques for managing data growth and costs.

SUMMARY

(3) Containerization of applications and open source deployments sometimes present special challenges in protecting data. Although they provide advantages in ease of deployment and/or reduced

costs, containerized and/or open source solutions are not necessarily equipped for robust and scalable data protection. For example, some database management systems (DBMS) ship without resources needed for data backups. For example, some open source backup utilities lack features for file level recovery, granular live browsing, and/or content indexing. These deficiencies require substantial manual intervention and/or tracking, thus contravening the aforementioned ease of deployment and cost advantages.

(4) The present inventors devised a streamlined approach that overcomes the deficiencies of the prior art. The illustrative computer-implemented technological improvement includes: (i) creating a specially-equipped “backup services container”, that comprises scripts implementing commands for accessing containerized applications, and also comprises execution resources and environments (enabling utilities) therefor, such as runtime C, python, etc.; (collectively the scripts and their enabling utilities are referred to herein as “backup preparation toolkits” or “backup toolkits” or “toolkits”); (ii) adding the backup services container to any number of Kubernetes pods comprising containerized applications (e.g., DBMS, web server, file system, login server, etc.); (iii) to provide maximum value and advantage, the backup services container is “over-equipped” with a spectrum of toolkits—generally many more toolkit types than types of target applications that might be found in any given Kubernetes pod; (iv) when a backup operation is pending for a target containerized application in the Kubernetes pod, the backup services container selects and applies a suitable backup preparation toolkit to the application to ready the application for the backup operation; (v) some embodiments include close interoperability with a proprietary data storage management system, thereby providing additional features and enhancements that are not possible with third-party backup systems; and (vi) the illustrative embodiments facilitate and enhance backup operations without necessitating changes to existing containers configured in the Kubernetes pod, thereby acting as value-added helpers that do not interfere with existing container configurations.

(5) Interoperability with an illustrative proprietary data storage management system provides additional advantages for managing backup operations and backed up data, such as creating inventories of containerized applications, creating and applying backup preferences for individual target applications, granular browse and restore features, indexing of backed up data, using retention policies, etc.

(6) However, in some embodiments, the illustrative backup services container does not require the proprietary data storage management system, and is capable of interoperating with third-party backup systems. A third-party backup system, such as Velero (see, e.g., <http://github.com/heptio/velero>), may include open source utilities for backing up Kubernetes applications, and generally lacks the robustness, feature richness, and scalability of the proprietary data storage management system described herein.

(7) Database management systems (DBMSs) and file systems typically require preparation before backup in order to maintain state and prevent data loss. Pre-backup preparation usually includes quiescing. Typically, a snapshot immediately follows to capture targeted data (e.g., database, file system data) while the application is quiescent, and the application is then released to resume normal operations. Other applications that need to preserve and maintain state require similar pre-backup preparation. The illustrative proprietary data storage management system, or a third-party backup system, then uses the snapshot as a data source to be processed for backup and to generate secondary copies therefrom.

(8) Because each application has its own unique architecture, features, states, and application programming interface (API), each application requires specialized pre-backup processing to prevent data loss and maintain state through the backup operation. Some applications are sometimes packaged without sufficient resources for implementing the necessary backup preparations. For example, a DBMS might lack suitable scripts for quiescing/unquiescing the DBMS. Even when scripts are included with the DBMS package, they need an execution environment (e.g., C runtime, python, etc.) that might be missing from the DBMS package. For example, applications supplied by an application store might lack necessary utilities/resources needed to enable scripts for accessing the applications. As a result, an attempt to back up the targeted application might fail altogether or might leave the

application in an inconsistent state resulting in data loss or data corruption.

(9) The illustrative backup services container is specially configured to overcome these risks and deficiencies by supplying a comprehensive set of backup preparation toolkits that include scripts and enabling execution environments. Each toolkit is suited to one or more types of applications. The collection of toolkits equipped into the backup services container thus has the ability to work with a variety of applications commonly deployed in containers and/or configured in Kubernetes pods. Since the backup services container is pre-configured before deployment into the Kubernetes pod, the more toolkits it comprises, the more versatile it will be in operation.

(10) Thus, the illustrative backup services container comprises a wide-ranging set of backup preparation toolkits targeted to a large variety of applications deployed in Kubernetes pods. In addition to the backup toolkits, the backup services container comprises logic for checking pod assets (e.g., applications, storage, etc.) and discovering their attributes; logic for selecting and applying suitable backup preparation toolkits to target applications and for releasing the applications after a backup operation completes; and logic for interfacing with external backup systems performing backup operations, such the illustrative proprietary data storage management system and/or third-party backup systems. Such backup systems are referred to herein as “external” in the sense that the backup systems are outside of the backup services container. The backup services container streamlines backup operations, and additionally functions as a collector of information that can be used productively by the proprietary data storage management system.

(11) In some embodiments, one or more components (e.g., data agents, media agents) of the proprietary data storage management system are configured within the backup services container to facilitate backup operations and/or to improve backup performance within the pod. In some embodiments, a Kubernetes pod is specially configured with components (e.g., data agents, media agents, storage manager, storage resources, etc.) of the proprietary data storage management system, thereby forming a “backup services pod.” The illustrative backup services pod interoperates with backup services container(s) in one or more other Kubernetes pods that co-reside within a Kubernetes node. The backup services pod also facilitates backup operations and/or improves backup performance for data in the Kubernetes node.

(12) Illustrative “discovery logic” in the backup services container determines what containerized applications are actually present in the Kubernetes pod and determines their characteristics and attributes. The discovery logic is configured to also interpret information, e.g., to determine whether “my_pictures” in a host container refers to a file system, a DBMS, a web server, or some other kind of application, by analyzing configuration parameters associated with “my_pictures” in the host container. In some embodiments, the discovery logic reports its findings, including inventories of applications and their attributes, to a storage manager in the proprietary data storage management system. The storage manager, which is generally responsible for controlling storage operations including backups, stores the information reported by the discovery logic. The storage manager further employs the received information to generate preferences that apply to the discovered container assets, e.g., storage policies, backup schedules, backup staggering, retention policies, etc. Certain information collected by the discovery logic is used in non-backup operations, e.g., mounting selected volumes in live browse, content indexing, etc.

(13) In some embodiments, the storage manager also establishes activity tracking of targeted containerized application to help trigger backup operations based on application activity rather than a fixed schedule (e.g., more frequent backups for busy applications, more frequent backups for applications generating large amounts of data, etc.). In some embodiments, the backup services container comprises activity monitoring logic (e.g., as part of the discovery logic and/or as a separate functional component) that tracks targeted containerized applications and reports to the external backup system, for example by using pre-defined thresholds to determine suitable timing for a backup.

(14) When a backup operation is triggered (e.g., by the proprietary data storage management system, by a third-party backup system, etc.), the backup services container receives notice, e.g., in the form of a trigger, an instruction, a message, etc. The discovery logic then determines and/or confirms which containerized applications (and other container assets) are present in the present Kubernetes pod,

including associated attributes.

(15) Logic for selecting and applying suitable backup preparation toolkits (“selection logic”) is then invoked at the backup services container. First, the selection logic determines which of the discovered containerized applications require which of the pre-configured backup preparation toolkits, if any. Some containerized applications need not be expressly prepared for backup, and therefore will not require the services of a backup toolkit. Next, the selection logic selects a suitable backup toolkit, establishes communications with the target containerized application, and executes the script in the toolkit, e.g., invoking certain commands via an API, thereby causing the target application to be readied for backup. For example, a backup preparation script issues one or more API commands to quiesce a DBMS or a file system. The selection logic reports to the external backup system that the target application is ready for backup (e.g., quiescent) and the selection logic waits for an indication that the backup operation has completed, e.g., received from the external backup system after a snapshot is taken. The selection logic then releases the application from its backup-ready (e.g., quiescent) state to resume normal operations (e.g., issues an unquiesce command via API). Any number of target containerized applications and different types and versions of containerized applications can be prepared for backup in this way using the various illustrative backup toolkits in the backup services container.

(16) The backup operation is typically managed by the external backup system, e.g., the proprietary data storage management system, a third-party backup system, etc. The backup services container uses an illustrative backup system interface logic to communicate with the external system for exchanging information, attributes, triggers, instructions, reports, status updates, etc. In some embodiments, the interface logic is part of the discovery and/or selection logic.

(17) The illustrative embodiments are directed to containerized applications configured in Kubernetes pods and/or Kubernetes nodes, and the illustrative containers are configured as Docker containers, but the invention is not so limited. In alternative embodiments, any operating system-level virtualization platform other than Docker containers and any container-orchestration system other than Kubernetes pods/nodes can be implemented with the illustrative backup services container and/or backup services pod. Moreover, although Kubernetes is often referred to in the context of serverless cloud computing environments, the invention is suitable for cloud and non-cloud implementations alike, without limitation.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1A is a block diagram illustrating an exemplary information management system.

(2) FIG. 1B is a detailed view of a primary storage device, a secondary storage device, and some examples of primary data and secondary copy data.

(3) FIG. 1C is a block diagram of an exemplary information management system including a storage manager, one or more data agents, and one or more media agents.

(4) FIG. 1D is a block diagram illustrating a scalable information management system.

(5) FIG. 1E illustrates certain secondary copy operations according to an exemplary storage policy.

(6) FIGS. 1F-1H are block diagrams illustrating suitable data structures that may be employed by the information management system.

(7) FIG. 2A illustrates a system and technique for synchronizing primary data to a destination such as a failover site using secondary copy data.

(8) FIG. 2B illustrates an information management system architecture incorporating use of a network file system (NFS) protocol for communicating between the primary and secondary storage subsystems.

(9) FIG. 2C is a block diagram of an example of a highly scalable managed data pool architecture.

(10) FIG. 3A is a block diagram illustrating some salient portions of a system 300 comprising backup services container 301 and data storage management system 302, according to an illustrative embodiment.

- (11) FIG. 3B is a block diagram illustrating backup services container **301** interoperating with a third-party backup system, according to another illustrative embodiment.
- (12) FIG. 3C is a block diagram of an illustrative overview of Kubernetes pods.
- (13) FIG. 3D is a block diagram of an illustrative overview of a Kubernetes node comprising a plurality of Kubernetes pods.
- (14) FIG. 4A is a block diagram depicting some salient details of backup services container **301** and data storage management system **302**.
- (15) FIG. 4B is a block diagram depicting some salient details of backup services container **301** and data storage management system **302**, wherein one or more data agents **442** are deployed in backup services container **301**.
- (16) FIG. 4C is a block diagram depicting some salient details of system **300** in which backup proxy **486** is configured and deployed as a Kubernetes pod within a Kubernetes node **413**.
- (17) FIG. 4D is a block diagram depicting some details of a “backup services pod” **486** deployed in a Kubernetes node.
- (18) FIG. 5 is a block diagram depicting some salient details of and logical information paths to/from backup services container **301** and storage manager **440**.
- (19) FIG. 6 depicts some salient operations of a method **600** according to an illustrative embodiment.
- (20) FIG. 7 depicts some salient details of block **608** in method **600**.
- (21) FIG. 8 depicts some salient details of block **610** in method **600**.

DETAILED DESCRIPTION

(22) Detailed descriptions and examples of systems and methods according to one or more illustrative embodiments of the present invention may be found in the section entitled PREPARING CONTAINERIZED APPLICATIONS FOR BACKUP USING A BACKUP SERVICES CONTAINER, as well as in the section entitled Example Embodiments, and also in FIGS. 3A-8 herein. Furthermore, components and functionality for preparing containerized applications for backup using a backup services container may be configured and/or incorporated into information management systems such as those described herein in FIGS. 1A-1H and 2A-2C.

(23) Various embodiments described herein are intimately tied to, enabled by, and would not exist except for, computer technology. For example, configuring a backup services container, communications to/from containerized applications and to/from external backup systems described herein in reference to various embodiments cannot reasonably be performed by humans alone, without the computer technology upon which they are implemented.

(24) Information Management System Overview

(25) With the increasing importance of protecting and leveraging data, organizations simply cannot risk losing critical data. Moreover, runaway data growth and other modern realities make protecting and managing data increasingly difficult. There is therefore a need for efficient, powerful, and user-friendly solutions for protecting and managing data and for smart and efficient management of data storage. Depending on the size of the organization, there may be many data production sources which are under the purview of tens, hundreds, or even thousands of individuals. In the past, individuals were sometimes responsible for managing and protecting their own data, and a patchwork of hardware and software point solutions may have been used in any given organization. These solutions were often provided by different vendors and had limited or no interoperability. Certain embodiments described herein address these and other shortcomings of prior approaches by implementing scalable, unified, organization-wide information management, including data storage management.

(26) FIG. 1A shows one such information management system **100** (or “system **100**”), which generally includes combinations of hardware and software configured to protect and manage data and metadata that are generated and used by computing devices in system **100**. System **100** may be referred to in some embodiments as a “storage management system” or a “data storage management system.” System **100** performs information management operations, some of which may be referred to as “storage operations” or “data storage operations,” to protect and manage the data residing in and/or managed by system **100**. The organization that employs system **100** may be a corporation or other business entity, non-profit organization, educational institution, household, governmental agency, or

the like.

(27) Generally, the systems and associated components described herein may be compatible with and/or provide some or all of the functionality of the systems and corresponding components described in one or more of the following U.S. patents/publications and patent applications assigned to Commvault Systems, Inc., each of which is hereby incorporated by reference in its entirety herein: U.S. Pat. No. 7,035,880, entitled “Modular Backup and Retrieval System Used in Conjunction With a Storage Area Network”; U.S. Pat. No. 7,107,298, entitled “System And Method For Archiving Objects In An Information Store”; U.S. Pat. No. 7,246,207, entitled “System and Method for Dynamically Performing Storage Operations in a Computer Network”; U.S. Pat. No. 7,315,923, entitled “System And Method For Combining Data Streams In Pipelined Storage Operations In A Storage Network”; U.S. Pat. No. 7,343,453, entitled “Hierarchical Systems and Methods for Providing a Unified View of Storage Information”; U.S. Pat. No. 7,395,282, entitled “Hierarchical Backup and Retrieval System”; U.S. Pat. No. 7,529,782, entitled “System and Methods for Performing a Snapshot and for Restoring Data”; U.S. Pat. No. 7,617,262, entitled “System and Methods for Monitoring Application Data in a Data Replication System”; U.S. Pat. No. 7,734,669, entitled “Managing Copies Of Data”; U.S. Pat. No. 7,747,579, entitled “Metabase for Facilitating Data Classification”; U.S. Pat. No. 8,156,086, entitled “Systems And Methods For Stored Data Verification”; U.S. Pat. No. 8,170,995, entitled “Method and System for Offline Indexing of Content and Classifying Stored Data”; U.S. Pat. No. 8,230,195, entitled “System And Method For Performing Auxiliary Storage Operations”; U.S. Pat. No. 8,285,681, entitled “Data Object Store and Server for a Cloud Storage Environment, Including Data Deduplication and Data Management Across Multiple Cloud Storage Sites”; U.S. Pat. No. 8,307,177, entitled “Systems And Methods For Management Of Virtualization Data”; U.S. Pat. No. 8,364,652, entitled “Content-Aligned, Block-Based Deduplication”; U.S. Pat. No. 8,578,120, entitled “Block-Level Single Instancing”; U.S. Pat. No. 8,954,446, entitled “Client-Side Repository in a Networked Deduplicated Storage System”; U.S. Pat. No. 9,020,900, entitled “Distributed Deduplicated Storage System”; U.S. Pat. No. 9,098,495, entitled “Application-Aware and Remote Single Instance Data Management”; U.S. Pat. No. 9,239,687, entitled “Systems and Methods for Retaining and Using Data Block Signatures in Data Protection Operations”; U.S. Pat. No. 9,633,033, entitled “High Availability Distributed Deduplicated Storage System”; U.S. Pat. Pub. No. 2006/0224846, entitled “System and Method to Support Single Instance Storage Operations”; U.S. Pat. Pub. No. 2016-0350391, entitled “Replication Using Deduplicated Secondary Copy Data”; U.S. Pat. Pub. No. 2017-0168903 A1, entitled “Live Synchronization and Management of Virtual Machines across Computing and Virtualization Platforms and Using Live Synchronization to Support Disaster Recovery”; U.S. Pat. Pub. No. 2017-0185488 A1, entitled “Application-Level Live Synchronization Across Computing Platforms Including Synchronizing Co-Resident Applications To Disparate Standby Destinations And Selectively Synchronizing Some Applications And Not Others”; U.S. Pat. Pub. No. 2017-0192866 A1, entitled “System For Redirecting Requests After A Secondary Storage Computing Device Failure”; U.S. Pat. Pub. No. 2017-0235647 A1, entitled “Data Protection Operations Based on Network Path Information”; and U.S. Pat. Pub. No. 2017-0242871 A1, entitled “Data Restoration Operations Based on Network Path Information”.

(28) System **100** includes computing devices and computing technologies. For instance, system **100** can include one or more client computing devices **102** and secondary storage computing devices **106**, as well as storage manager **140** or a host computing device for it. Computing devices can include, without limitation, one or more: workstations, personal computers, desktop computers, or other types of generally fixed computing systems such as mainframe computers, servers, and minicomputers. Other computing devices can include mobile or portable computing devices, such as one or more laptops, tablet computers, personal data assistants, mobile phones (such as smartphones), and other mobile or portable computing devices such as embedded computers, set top boxes, vehicle-mounted devices, wearable computers, etc. Servers can include mail servers, file servers, database servers, virtual machine servers, and web servers. Any given computing device comprises one or more processors (e.g., CPU and/or single-core or multi-core processors), as well as corresponding non-transitory computer memory (e.g., random-access memory (RAM)) for storing computer programs

which are to be executed by the one or more processors. Other computer memory for mass storage of data may be packaged/configured with the computing device (e.g., an internal hard disk) and/or may be external and accessible by the computing device (e.g., network-attached storage, a storage array, etc.). In some cases, a computing device includes cloud computing resources, which may be implemented as virtual machines. For instance, one or more virtual machines may be provided to the organization by a third-party cloud service vendor.

(29) In some embodiments, computing devices can include one or more virtual machine(s) running on a physical host computing device (or “host machine”) operated by the organization. As one example, the organization may use one virtual machine as a database server and another virtual machine as a mail server, both virtual machines operating on the same host machine. A Virtual machine (“VM”) is a software implementation of a computer that does not physically exist and is instead instantiated in an operating system of a physical computer (or host machine) to enable applications to execute within the VM's environment, i.e., a VM emulates a physical computer. A VM includes an operating system and associated virtual resources, such as computer memory and processor(s). A hypervisor operates between the VM and the hardware of the physical host machine and is generally responsible for creating and running the VMs. Hypervisors are also known in the art as virtual machine monitors or a virtual machine managers or “VMMs”, and may be implemented in software, firmware, and/or specialized hardware installed on the host machine. Examples of hypervisors include ESX Server, by VMware, Inc. of Palo Alto, California; Microsoft Virtual Server and Microsoft Windows Server Hyper-V, both by Microsoft Corporation of Redmond, Washington; Sun xVM by Oracle America Inc. of Santa Clara, California; and Xen by Citrix Systems, Santa Clara, California. The hypervisor provides resources to each virtual operating system such as a virtual processor, virtual memory, a virtual network device, and a virtual disk. Each virtual machine has one or more associated virtual disks. The hypervisor typically stores the data of virtual disks in files on the file system of the physical host machine, called virtual machine disk files (“VMDK” in VMware lingo) or virtual hard disk image files (in Microsoft lingo). For example, VMware's ESX Server provides the Virtual Machine File System (VMFS) for the storage of virtual machine disk files. A virtual machine reads data from and writes data to its virtual disk much the way that a physical machine reads data from and writes data to a physical disk. Examples of techniques for implementing information management in a cloud computing environment are described in U.S. Pat. No. 8,285,681. Examples of techniques for implementing information management in a virtualized computing environment are described in U.S. Pat. No. 8,307,177.

(30) Information management system **100** can also include electronic data storage devices, generally used for mass storage of data, including, e.g., primary storage devices **104** and secondary storage devices **108**. Storage devices can generally be of any suitable type including, without limitation, disk drives, storage arrays (e.g., storage-area network (SAN) and/or network-attached storage (NAS) technology), semiconductor memory (e.g., solid state storage devices), network attached storage (NAS) devices, tape libraries, or other magnetic, non-tape storage devices, optical media storage devices, combinations of the same, etc. In some embodiments, storage devices form part of a distributed file system. In some cases, storage devices are provided in a cloud storage environment (e.g., a private cloud or one operated by a third-party vendor), whether for primary data or secondary copies or both.

(31) Depending on context, the term “information management system” can refer to generally all of the illustrated hardware and software components in FIG. 1C, or the term may refer to only a subset of the illustrated components. For instance, in some cases, system **100** generally refers to a combination of specialized components used to protect, move, manage, manipulate, analyze, and/or process data and metadata generated by client computing devices **102**. However, system **100** in some cases does not include the underlying components that generate and/or store primary data **112**, such as the client computing devices **102** themselves, and the primary storage devices **104**. Likewise secondary storage devices **108** (e.g., a third-party provided cloud storage environment) may not be part of system **100**. As an example, “information management system” or “storage management system” may sometimes refer to one or more of the following components, which will be described in further detail below: storage

manager, data agent, and media agent.

(32) One or more client computing devices **102** may be part of system **100**, each client computing device **102** having an operating system and at least one application **110** and one or more accompanying data agents executing thereon; and associated with one or more primary storage devices **104** storing primary data **112**. Client computing device(s) **102** and primary storage devices **104** may generally be referred to in some cases as primary storage subsystem **117**.

(33) Client Computing Devices, Clients, and Subclients

(34) Typically, a variety of sources in an organization produce data to be protected and managed. As just one illustrative example, in a corporate environment such data sources can be employee workstations and company servers such as a mail server, a web server, a database server, a transaction server, or the like. In system **100**, data generation sources include one or more client computing devices **102**. A computing device that has a data agent **142** installed and operating on it is generally referred to as a “client computing device” **102**, and may include any type of computing device, without limitation. A client computing device **102** may be associated with one or more users and/or user accounts.

(35) A “client” is a logical component of information management system **100**, which may represent a logical grouping of one or more data agents installed on a client computing device **102**. Storage manager **140** recognizes a client as a component of system **100**, and in some embodiments, may automatically create a client component the first time a data agent **142** is installed on a client computing device **102**. Because data generated by executable component(s) **110** is tracked by the associated data agent **142** so that it may be properly protected in system **100**, a client may be said to generate data and to store the generated data to primary storage, such as primary storage device **104**. However, the terms “client” and “client computing device” as used herein do not imply that a client computing device **102** is necessarily configured in the client/server sense relative to another computing device such as a mail server, or that a client computing device **102** cannot be a server in its own right. As just a few examples, a client computing device **102** can be and/or include mail servers, file servers, database servers, virtual machine servers, and/or web servers.

(36) Each client computing device **102** may have application(s) **110** executing thereon which generate and manipulate the data that is to be protected from loss and managed in system **100**. Applications **110** generally facilitate the operations of an organization, and can include, without limitation, mail server applications (e.g., Microsoft Exchange Server), file system applications, mail client applications (e.g., Microsoft Exchange Client), database applications or database management systems (e.g., SQL, Oracle, SAP, Lotus Notes Database), word processing applications (e.g., Microsoft Word), spreadsheet applications, financial applications, presentation applications, graphics and/or video applications, browser applications, mobile applications, entertainment applications, and so on. Each application **110** may be accompanied by an application-specific data agent **142**, though not all data agents **142** are application-specific or associated with only application. A file manager application, e.g., Microsoft Windows Explorer, may be considered an application **110** and may be accompanied by its own data agent **142**. Client computing devices **102** can have at least one operating system (e.g., Microsoft Windows, Mac OS X, iOS, IBM z/OS, Linux, other Unix-based operating systems, etc.) installed thereon, which may support or host one or more file systems and other applications **110**. In some embodiments, a virtual machine that executes on a host client computing device **102** may be considered an application **110** and may be accompanied by a specific data agent **142** (e.g., virtual server data agent).

(37) Client computing devices **102** and other components in system **100** can be connected to one another via one or more electronic communication pathways **114**. For example, a first communication pathway **114** may communicatively couple client computing device **102** and secondary storage computing device **106**; a second communication pathway **114** may communicatively couple storage manager **140** and client computing device **102**; and a third communication pathway **114** may communicatively couple storage manager **140** and secondary storage computing device **106**, etc. (see, e.g., FIG. 1A and FIG. 1C). A communication pathway **114** can include one or more networks or other connection types including one or more of the following, without limitation: the Internet, a wide area

network (WAN), a local area network (LAN), a Storage Area Network (SAN), a Fibre Channel (FC) connection, a Small Computer System Interface (SCSI) connection, a virtual private network (VPN), a token ring or TCP/IP based network, an intranet network, a point-to-point link, a cellular network, a wireless data transmission system, a two-way cable system, an interactive kiosk network, a satellite network, a broadband network, a baseband network, a neural network, a mesh network, an ad hoc network, other appropriate computer or telecommunications networks, combinations of the same or the like. Communication pathways **114** in some cases may also include application programming interfaces (APIs) including, e.g., cloud service provider APIs, virtual machine management APIs, and hosted service provider APIs. The underlying infrastructure of communication pathways **114** may be wired and/or wireless, analog and/or digital, or any combination thereof; and the facilities used may be private, public, third-party provided, or any combination thereof, without limitation.

(38) A “subclient” is a logical grouping of all or part of a client's primary data **112**. In general, a subclient may be defined according to how the subclient data is to be protected as a unit in system **100**. For example, a subclient may be associated with a certain storage policy. A given client may thus comprise several subclients, each subclient associated with a different storage policy. For example, some files may form a first subclient that requires compression and deduplication and is associated with a first storage policy. Other files of the client may form a second subclient that requires a different retention schedule as well as encryption, and may be associated with a different, second storage policy. As a result, though the primary data may be generated by the same application **110** and may belong to one given client, portions of the data may be assigned to different subclients for distinct treatment by system **100**. More detail on subclients is given in regard to storage policies below.

(39) Primary Data and Exemplary Primary Storage Devices

(40) Primary data **112** is generally production data or “live” data generated by the operating system and/or applications **110** executing on client computing device **102**. Primary data **112** is generally stored on primary storage device(s) **104** and is organized via a file system operating on the client computing device **102**. Thus, client computing device(s) **102** and corresponding applications **110** may create, access, modify, write, delete, and otherwise use primary data **112**. Primary data **112** is generally in the native format of the source application **110**. Primary data **112** is an initial or first stored body of data generated by the source application **110**. Primary data **112** in some cases is created substantially directly from data generated by the corresponding source application **110**. It can be useful in performing certain tasks to organize primary data **112** into units of different granularities. In general, primary data **112** can include files, directories, file system volumes, data blocks, extents, or any other hierarchies or organizations of data objects. As used herein, a “data object” can refer to (i) any file that is currently addressable by a file system or that was previously addressable by the file system (e.g., an archive file), and/or to (ii) a subset of such a file (e.g., a data block, an extent, etc.). Primary data **112** may include structured data (e.g., database files), unstructured data (e.g., documents), and/or semi-structured data. See, e.g., FIG. **1B**.

(41) It can also be useful in performing certain functions of system **100** to access and modify metadata within primary data **112**. Metadata generally includes information about data objects and/or characteristics associated with the data objects. For simplicity herein, it is to be understood that, unless expressly stated otherwise, any reference to primary data **112** generally also includes its associated metadata, but references to metadata generally do not include the primary data. Metadata can include, without limitation, one or more of the following: the data owner (e.g., the client or user that generates the data), the last modified time (e.g., the time of the most recent modification of the data object), a data object name (e.g., a file name), a data object size (e.g., a number of bytes of data), information about the content (e.g., an indication as to the existence of a particular search term), user-supplied tags, to/from information for email (e.g., an email sender, recipient, etc.), creation date, file type (e.g., format or application type), last accessed time, application type (e.g., type of application that generated the data object), location/network (e.g., a current, past or future location of the data object and network pathways to/from the data object), geographic location (e.g., GPS coordinates), frequency of change (e.g., a period in which the data object is modified), business unit (e.g., a group or department that generates, manages or is otherwise associated with the data object), aging information (e.g., a

schedule, such as a time period, in which the data object is migrated to secondary or long term storage), boot sectors, partition layouts, file location within a file folder directory structure, user permissions, owners, groups, access control lists (ACLs), system metadata (e.g., registry information), combinations of the same or other similar information related to the data object. In addition to metadata generated by or related to file systems and operating systems, some applications **110** and/or other components of system **100** maintain indices of metadata for data objects, e.g., metadata associated with individual email messages. The use of metadata to perform classification and other functions is described in greater detail below.

(42) Primary storage devices **104** storing primary data **112** may be relatively fast and/or expensive technology (e.g., flash storage, a disk drive, a hard-disk storage array, solid state memory, etc.), typically to support high-performance live production environments. Primary data **112** may be highly changeable and/or may be intended for relatively short term retention (e.g., hours, days, or weeks). According to some embodiments, client computing device **102** can access primary data **112** stored in primary storage device **104** by making conventional file system calls via the operating system. Each client computing device **102** is generally associated with and/or in communication with one or more primary storage devices **104** storing corresponding primary data **112**. A client computing device **102** is said to be associated with or in communication with a particular primary storage device **104** if it is capable of one or more of: routing and/or storing data (e.g., primary data **112**) to the primary storage device **104**, coordinating the routing and/or storing of data to the primary storage device **104**, retrieving data from the primary storage device **104**, coordinating the retrieval of data from the primary storage device **104**, and modifying and/or deleting data in the primary storage device **104**. Thus, a client computing device **102** may be said to access data stored in an associated storage device **104**.

(43) Primary storage device **104** may be dedicated or shared. In some cases, each primary storage device **104** is dedicated to an associated client computing device **102**, e.g., a local disk drive. In other cases, one or more primary storage devices **104** can be shared by multiple client computing devices **102**, e.g., via a local network, in a cloud storage implementation, etc. As one example, primary storage device **104** can be a storage array shared by a group of client computing devices **102**, such as EMC Clariion, EMC Symmetrix, EMC Celerra, Dell EqualLogic, IBM XIV, NetApp FAS, HP EVA, and HP 3PAR.

(44) System **100** may also include hosted services (not shown), which may be hosted in some cases by an entity other than the organization that employs the other components of system **100**. For instance, the hosted services may be provided by online service providers. Such service providers can provide social networking services, hosted email services, or hosted productivity applications or other hosted applications such as software-as-a-service (SaaS), platform-as-a-service (PaaS), application service providers (ASPs), cloud services, or other mechanisms for delivering functionality via a network. As it services users, each hosted service may generate additional data and metadata, which may be managed by system **100**, e.g., as primary data **112**. In some cases, the hosted services may be accessed using one of the applications **110**. As an example, a hosted mail service may be accessed via browser running on a client computing device **102**.

(45) Secondary Copies and Exemplary Secondary Storage Devices

(46) Primary data **112** stored on primary storage devices **104** may be compromised in some cases, such as when an employee deliberately or accidentally deletes or overwrites primary data **112**. Or primary storage devices **104** can be damaged, lost, or otherwise corrupted. For recovery and/or regulatory compliance purposes, it is therefore useful to generate and maintain copies of primary data **112**.

Accordingly, system **100** includes one or more secondary storage computing devices **106** and one or more secondary storage devices **108** configured to create and store one or more secondary copies **116** of primary data **112** including its associated metadata. The secondary storage computing devices **106** and the secondary storage devices **108** may be referred to as secondary storage subsystem **118**.

(47) Secondary copies **116** can help in search and analysis efforts and meet other information management goals as well, such as: restoring data and/or metadata if an original version is lost (e.g., by deletion, corruption, or disaster); allowing point-in-time recovery; complying with regulatory data

retention and electronic discovery (e-discovery) requirements; reducing utilized storage capacity in the production system and/or in secondary storage; facilitating organization and search of data; improving user access to data files across multiple computing devices and/or hosted services; and implementing data retention and pruning policies.

(48) A secondary copy **116** can comprise a separate stored copy of data that is derived from one or more earlier-created stored copies (e.g., derived from primary data **112** or from another secondary copy **116**). Secondary copies **116** can include point-in-time data, and may be intended for relatively long-term retention before some or all of the data is moved to other storage or discarded. In some cases, a secondary copy **116** may be in a different storage device than other previously stored copies; and/or may be remote from other previously stored copies. Secondary copies **116** can be stored in the same storage device as primary data **112**. For example, a disk array capable of performing hardware snapshots stores primary data **112** and creates and stores hardware snapshots of the primary data **112** as secondary copies **116**. Secondary copies **116** may be stored in relatively slow and/or lower cost storage (e.g., magnetic tape). A secondary copy **116** may be stored in a backup or archive format, or in some other format different from the native source application format or other format of primary data **112**.

(49) Secondary storage computing devices **106** may index secondary copies **116** (e.g., using a media agent **144**), enabling users to browse and restore at a later time and further enabling the lifecycle management of the indexed data. After creation of a secondary copy **116** that represents certain primary data **112**, a pointer or other location indicia (e.g., a stub) may be placed in primary data **112**, or be otherwise associated with primary data **112**, to indicate the current location of a particular secondary copy **116**. Since an instance of a data object or metadata in primary data **112** may change over time as it is modified by application **110** (or hosted service or the operating system), system **100** may create and manage multiple secondary copies **116** of a particular data object or metadata, each copy representing the state of the data object in primary data **112** at a particular point in time.

Moreover, since an instance of a data object in primary data **112** may eventually be deleted from primary storage device **104** and the file system, system **100** may continue to manage point-in-time representations of that data object, even though the instance in primary data **112** no longer exists. For virtual machines, the operating system and other applications **110** of client computing device(s) **102** may execute within or under the management of virtualization software (e.g., a VMM), and the primary storage device(s) **104** may comprise a virtual disk created on a physical storage device. System **100** may create secondary copies **116** of the files or other data objects in a virtual disk file and/or secondary copies **116** of the entire virtual disk file itself (e.g., of an entire .vmdk file).

(50) Secondary copies **116** are distinguishable from corresponding primary data **112**. First, secondary copies **116** can be stored in a different format from primary data **112** (e.g., backup, archive, or other non-native format). For this or other reasons, secondary copies **116** may not be directly usable by applications **110** or client computing device **102** (e.g., via standard system calls or otherwise) without modification, processing, or other intervention by system **100** which may be referred to as “restore” operations. Secondary copies **116** may have been processed by data agent **142** and/or media agent **144** in the course of being created (e.g., compression, deduplication, encryption, integrity markers, indexing, formatting, application-aware metadata, etc.), and thus secondary copy **116** may represent source primary data **112** without necessarily being exactly identical to the source.

(51) Second, secondary copies **116** may be stored on a secondary storage device **108** that is inaccessible to application **110** running on client computing device **102** and/or hosted service. Some secondary copies **116** may be “offline copies,” in that they are not readily available (e.g., not mounted to tape or disk). Offline copies can include copies of data that system **100** can access without human intervention (e.g., tapes within an automated tape library, but not yet mounted in a drive), and copies that the system **100** can access only with some human intervention (e.g., tapes located at an offsite storage site).

(52) Using Intermediate Devices for Creating Secondary Copies—Secondary Storage Computing Devices

(53) Creating secondary copies can be challenging when hundreds or thousands of client computing devices **102** continually generate large volumes of primary data **112** to be protected. Also, there can be

significant overhead involved in the creation of secondary copies **116**. Moreover, specialized programmed intelligence and/or hardware capability is generally needed for accessing and interacting with secondary storage devices **108**. Client computing devices **102** may interact directly with a secondary storage device **108** to create secondary copies **116**, but in view of the factors described above, this approach can negatively impact the ability of client computing device **102** to serve/service application **110** and produce primary data **112**. Further, any given client computing device **102** may not be optimized for interaction with certain secondary storage devices **108**.

(54) Thus, system **100** may include one or more software and/or hardware components which generally act as intermediaries between client computing devices **102** (that generate primary data **112**) and secondary storage devices **108** (that store secondary copies **116**). In addition to off-loading certain responsibilities from client computing devices **102**, these intermediate components provide other benefits. For instance, as discussed further below with respect to FIG. **1D**, distributing some of the work involved in creating secondary copies **116** can enhance scalability and improve system performance. For instance, using specialized secondary storage computing devices **106** and media agents **144** for interfacing with secondary storage devices **108** and/or for performing certain data processing operations can greatly improve the speed with which system **100** performs information management operations and can also improve the capacity of the system to handle large numbers of such operations, while reducing the computational load on the production environment of client computing devices **102**. The intermediate components can include one or more secondary storage computing devices **106** as shown in FIG. **1A** and/or one or more media agents **144**. Media agents are discussed further below (e.g., with respect to FIGS. **1C-1E**). These special-purpose components of system **100** comprise specialized programmed intelligence and/or hardware capability for writing to, reading from, instructing, communicating with, or otherwise interacting with secondary storage devices **108**.

(55) Secondary storage computing device(s) **106** can comprise any of the computing devices described above, without limitation. In some cases, secondary storage computing device(s) **106** also include specialized hardware componentry and/or software intelligence (e.g., specialized interfaces) for interacting with certain secondary storage device(s) **108** with which they may be specially associated.

(56) To create a secondary copy **116** involving the copying of data from primary storage subsystem **117** to secondary storage subsystem **118**, client computing device **102** may communicate the primary data **112** to be copied (or a processed version thereof generated by a data agent **142**) to the designated secondary storage computing device **106**, via a communication pathway **114**. Secondary storage computing device **106** in turn may further process and convey the data or a processed version thereof to secondary storage device **108**. One or more secondary copies **116** may be created from existing secondary copies **116**, such as in the case of an auxiliary copy operation, described further below.

(57) Exemplary Primary Data and an Exemplary Secondary Copy

(58) FIG. **1B** is a detailed view of some specific examples of primary data stored on primary storage device(s) **104** and secondary copy data stored on secondary storage device(s) **108**, with other components of the system removed for the purposes of illustration. Stored on primary storage device(s) **104** are primary data **112** objects including word processing documents **119A-B**, spreadsheets **120**, presentation documents **122**, video files **124**, image files **126**, email mailboxes **128** (and corresponding email messages **129A-C**), HTML/XML or other types of markup language files **130**, databases **132** and corresponding tables or other data structures **133A-133C**. Some or all primary data **112** objects are associated with corresponding metadata (e.g., "Meta1-11"), which may include file system metadata and/or application-specific metadata. Stored on the secondary storage device(s) **108** are secondary copy **116** data objects **134A-C** which may include copies of or may otherwise represent corresponding primary data **112**.

(59) Secondary copy data objects **134A-C** can individually represent more than one primary data object. For example, secondary copy data object **134A** represents three separate primary data objects **133C**, **122**, and **129C** (represented as **133C'**, **122'**, and **129C'**, respectively, and accompanied by corresponding metadata Meta11, Meta3, and Meta8, respectively). Moreover, as indicated by the prime mark ('), secondary storage computing devices **106** or other components in secondary storage

subsystem **118** may process the data received from primary storage subsystem **117** and store a secondary copy including a transformed and/or supplemented representation of a primary data object and/or metadata that is different from the original format, e.g., in a compressed, encrypted, deduplicated, or other modified format. For instance, secondary storage computing devices **106** can generate new metadata or other information based on said processing, and store the newly generated information along with the secondary copies. Secondary copy data object **134B** represents primary data objects **120**, **133B**, and **119A** as **120'**, **133B'**, and **119A'**, respectively, accompanied by corresponding metadata Meta2, Meta10, and Meta1, respectively. Also, secondary copy data object **134C** represents primary data objects **133A**, **119B**, and **129A** as **133A'**, **119B'**, and **129A'**, respectively, accompanied by corresponding metadata Meta9, Meta5, and Meta6, respectively.

(60) Exemplary Information Management System Architecture

(61) System **100** can incorporate a variety of different hardware and software components, which can in turn be organized with respect to one another in many different configurations, depending on the embodiment. There are critical design choices involved in specifying the functional responsibilities of the components and the role of each component in system **100**. Such design choices can impact how system **100** performs and adapts to data growth and other changing circumstances. FIG. **1C** shows a system **100** designed according to these considerations and includes: storage manager **140**, one or more data agents **142** executing on client computing device(s) **102** and configured to process primary data **112**, and one or more media agents **144** executing on one or more secondary storage computing devices **106** for performing tasks involving secondary storage devices **108**.

(62) Storage Manager

(63) Storage manager **140** is a centralized storage and/or information manager that is configured to perform certain control functions and also to store certain critical information about system **100**—hence storage manager **140** is said to manage system **100**. As noted, the number of components in system **100** and the amount of data under management can be large. Managing the components and data is therefore a significant task, which can grow unpredictably as the number of components and data scale to meet the needs of the organization. For these and other reasons, according to certain embodiments, responsibility for controlling system **100**, or at least a significant portion of that responsibility, is allocated to storage manager **140**. Storage manager **140** can be adapted independently according to changing circumstances, without having to replace or re-design the remainder of the system. Moreover, a computing device for hosting and/or operating as storage manager **140** can be selected to best suit the functions and networking needs of storage manager **140**. These and other advantages are described in further detail below and with respect to FIG. **1D**.

(64) Storage manager **140** may be a software module or other application hosted by a suitable computing device. In some embodiments, storage manager **140** is itself a computing device that performs the functions described herein. Storage manager **140** comprises or operates in conjunction with one or more associated data structures such as a dedicated database (e.g., management database **146**), depending on the configuration. The storage manager **140** generally initiates, performs, coordinates, and/or controls storage and other information management operations performed by system **100**, e.g., to protect and control primary data **112** and secondary copies **116**. In general, storage manager **140** is said to manage system **100**, which includes communicating with, instructing, and controlling in some circumstances components such as data agents **142** and media agents **144**, etc.

(65) As shown by the dashed arrowed lines **114** in FIG. **1C**, storage manager **140** may communicate with, instruct, and/or control some or all elements of system **100**, such as data agents **142** and media agents **144**. In this manner, storage manager **140** manages the operation of various hardware and software components in system **100**. In certain embodiments, control information originates from storage manager **140** and status as well as index reporting is transmitted to storage manager **140** by the managed components, whereas payload data and metadata are generally communicated between data agents **142** and media agents **144** (or otherwise between client computing device(s) **102** and secondary storage computing device(s) **106**), e.g., at the direction of and under the management of storage manager **140**. Control information can generally include parameters and instructions for carrying out information management operations, such as, without limitation, instructions to perform a task

associated with an operation, timing information specifying when to initiate a task, data path information specifying what components to communicate with or access in carrying out an operation, and the like. In other embodiments, some information management operations are controlled or initiated by other components of system **100** (e.g., by media agents **144** or data agents **142**), instead of or in combination with storage manager **140**.

(66) According to certain embodiments, storage manager **140** provides one or more of the following functions: communicating with data agents **142** and media agents **144**, including transmitting instructions, messages, and/or queries, as well as receiving status reports, index information, messages, and/or queries, and responding to same; initiating execution of information management operations; initiating restore and recovery operations; managing secondary storage devices **108** and inventory/capacity of the same; allocating secondary storage devices **108** for secondary copy operations; reporting, searching, and/or classification of data in system **100**; monitoring completion of and status reporting related to information management operations and jobs; tracking movement of data within system **100**; tracking age information relating to secondary copies **116**, secondary storage devices **108**, comparing the age information against retention guidelines, and initiating data pruning when appropriate; tracking logical associations between components in system **100**; protecting metadata associated with system **100**, e.g., in management database **146**; implementing job management, schedule management, event management, alert management, reporting, job history maintenance, user security management, disaster recovery management, and/or user interfacing for system administrators and/or end users of system **100**; sending, searching, and/or viewing of log files; and implementing operations management functionality.

(67) Storage manager **140** may maintain an associated database **146** (or “storage manager database **146**” or “management database **146**”) of management-related data and information management policies **148**. Database **146** is stored in computer memory accessible by storage manager **140**. Database **146** may include a management index **150** (or “index **150**”) or other data structure(s) that may store: logical associations between components of the system; user preferences and/or profiles (e.g., preferences regarding encryption, compression, or deduplication of primary data or secondary copies; preferences regarding the scheduling, type, or other aspects of secondary copy or other operations; mappings of particular information management users or user accounts to certain computing devices or other components, etc.; management tasks; media containerization; other useful data; and/or any combination thereof. For example, storage manager **140** may use index **150** to track logical associations between media agents **144** and secondary storage devices **108** and/or movement of data to/from secondary storage devices **108**. For instance, index **150** may store data associating a client computing device **102** with a particular media agent **144** and/or secondary storage device **108**, as specified in an information management policy **148**.

(68) Administrators and others may configure and initiate certain information management operations on an individual basis. But while this may be acceptable for some recovery operations or other infrequent tasks, it is often not workable for implementing on-going organization-wide data protection and management. Thus, system **100** may utilize information management policies **148** for specifying and executing information management operations on an automated basis. Generally, an information management policy **148** can include a stored data structure or other information source that specifies parameters (e.g., criteria and rules) associated with storage management or other information management operations. Storage manager **140** can process an information management policy **148** and/or index **150** and, based on the results, identify an information management operation to perform, identify the appropriate components in system **100** to be involved in the operation (e.g., client computing devices **102** and corresponding data agents **142**, secondary storage computing devices **106** and corresponding media agents **144**, etc.), establish connections to those components and/or between those components, and/or instruct and control those components to carry out the operation. In this manner, system **100** can translate stored information into coordinated activity among the various computing devices in system **100**.

(69) Management database **146** may maintain information management policies **148** and associated data, although information management policies **148** can be stored in computer memory at any

appropriate location outside management database **146**. For instance, an information management policy **148** such as a storage policy may be stored as metadata in a media agent database **152** or in a secondary storage device **108** (e.g., as an archive copy) for use in restore or other information management operations, depending on the embodiment. Information management policies **148** are described further below. According to certain embodiments, management database **146** comprises a relational database (e.g., an SQL database) for tracking metadata, such as metadata associated with secondary copy operations (e.g., what client computing devices **102** and corresponding subclient data were protected and where the secondary copies are stored and which media agent **144** performed the storage operation(s)). This and other metadata may additionally be stored in other locations, such as at secondary storage computing device **106** or on the secondary storage device **108**, allowing data recovery without the use of storage manager **140** in some cases. Thus, management database **146** may comprise data needed to kick off secondary copy operations (e.g., storage policies, schedule policies, etc.), status and reporting information about completed jobs (e.g., status and error reports on yesterday's backup jobs), and additional information sufficient to enable restore and disaster recovery operations (e.g., media agent associations, location indexing, content indexing, etc.).

(70) Storage manager **140** may include a jobs agent **156**, a user interface **158**, and a management agent **154**, all of which may be implemented as interconnected software modules or application programs. These are described further below.

(71) Jobs agent **156** in some embodiments initiates, controls, and/or monitors the status of some or all information management operations previously performed, currently being performed, or scheduled to be performed by system **100**. A job is a logical grouping of information management operations such as daily storage operations scheduled for a certain set of subclients (e.g., generating incremental block-level backup copies **116** at a certain time every day for database files in a certain geographical location). Thus, jobs agent **156** may access information management policies **148** (e.g., in management database **146**) to determine when, where, and how to initiate/control jobs in system **100**.

(72) Storage Manager User Interfaces

(73) User interface **158** may include information processing and display software, such as a graphical user interface (GUI), an application program interface (API), and/or other interactive interface(s) through which users and system processes can retrieve information about the status of information management operations or issue instructions to storage manager **140** and other components. Via user interface **158**, users may issue instructions to the components in system **100** regarding performance of secondary copy and recovery operations. For example, a user may modify a schedule concerning the number of pending secondary copy operations. As another example, a user may employ the GUI to view the status of pending secondary copy jobs or to monitor the status of certain components in system **100** (e.g., the amount of capacity left in a storage device). Storage manager **140** may track information that permits it to select, designate, or otherwise identify content indices, deduplication databases, or similar databases or resources or data sets within its information management cell (or another cell) to be searched in response to certain queries. Such queries may be entered by the user by interacting with user interface **158**.

(74) Various embodiments of information management system **100** may be configured and/or designed to generate user interface data usable for rendering the various interactive user interfaces described. The user interface data may be used by system **100** and/or by another system, device, and/or software program (for example, a browser program), to render the interactive user interfaces. The interactive user interfaces may be displayed on, for example, electronic displays (including, for example, touch-enabled displays), consoles, etc., whether direct-connected to storage manager **140** or communicatively coupled remotely, e.g., via an internet connection. The present disclosure describes various embodiments of interactive and dynamic user interfaces, some of which may be generated by user interface agent **158**, and which are the result of significant technological development. The user interfaces described herein may provide improved human-computer interactions, allowing for significant cognitive and ergonomic efficiencies and advantages over previous systems, including reduced mental workloads, improved decision-making, and the like. User interface **158** may operate in a single integrated view or console (not shown). The console may support a reporting capability for

generating a variety of reports, which may be tailored to a particular aspect of information management.

(75) User interfaces are not exclusive to storage manager **140** and in some embodiments a user may access information locally from a computing device component of system **100**. For example, some information pertaining to installed data agents **142** and associated data streams may be available from client computing device **102**. Likewise, some information pertaining to media agents **144** and associated data streams may be available from secondary storage computing device **106**.

(76) Storage Manager Management Agent

(77) Management agent **154** can provide storage manager **140** with the ability to communicate with other components within system **100** and/or with other information management cells via network protocols and application programming interfaces (APIs) including, e.g., HTTP, HTTP, FTP, REST, virtualization software APIs, cloud service provider APIs, and hosted service provider APIs, without limitation. Management agent **154** also allows multiple information management cells to communicate with one another. For example, system **100** in some cases may be one information management cell in a network of multiple cells adjacent to one another or otherwise logically related, e.g., in a WAN or LAN. With this arrangement, the cells may communicate with one another through respective management agents **154**. Inter-cell communications and hierarchy is described in greater detail in e.g., U.S. Pat. No. 7,343,453.

(78) Information Management Cell

(79) An “information management cell” (or “storage operation cell” or “cell”) may generally include a logical and/or physical grouping of a combination of hardware and software components associated with performing information management operations on electronic data, typically one storage manager **140** and at least one data agent **142** (executing on a client computing device **102**) and at least one media agent **144** (executing on a secondary storage computing device **106**). For instance, the components shown in FIG. **1C** may together form an information management cell. Thus, in some configurations, a system **100** may be referred to as an information management cell or a storage operation cell. A given cell may be identified by the identity of its storage manager **140**, which is generally responsible for managing the cell.

(80) Multiple cells may be organized hierarchically, so that cells may inherit properties from hierarchically superior cells or be controlled by other cells in the hierarchy (automatically or otherwise). Alternatively, in some embodiments, cells may inherit or otherwise be associated with information management policies, preferences, information management operational parameters, or other properties or characteristics according to their relative position in a hierarchy of cells. Cells may also be organized hierarchically according to function, geography, architectural considerations, or other factors useful or desirable in performing information management operations. For example, a first cell may represent a geographic segment of an enterprise, such as a Chicago office, and a second cell may represent a different geographic segment, such as a New York City office. Other cells may represent departments within a particular office, e.g., human resources, finance, engineering, etc. Where delineated by function, a first cell may perform one or more first types of information management operations (e.g., one or more first types of secondary copies at a certain frequency), and a second cell may perform one or more second types of information management operations (e.g., one or more second types of secondary copies at a different frequency and under different retention rules). In general, the hierarchical information is maintained by one or more storage managers **140** that manage the respective cells (e.g., in corresponding management database(s) **146**).

(81) Data Agents

(82) A variety of different applications **110** can operate on a given client computing device **102**, including operating systems, file systems, database applications, e-mail applications, and virtual machines, just to name a few. And, as part of the process of creating and restoring secondary copies **116**, the client computing device **102** may be tasked with processing and preparing the primary data **112** generated by these various applications **110**. Moreover, the nature of the processing/preparation can differ across application types, e.g., due to inherent structural, state, and formatting differences among applications **110** and/or the operating system of client computing device **102**. Each data agent

142 is therefore advantageously configured in some embodiments to assist in the performance of information management operations based on the type of data that is being protected at a client-specific and/or application-specific level.

(83) Data agent **142** is a component of information system **100** and is generally directed by storage manager **140** to participate in creating or restoring secondary copies **116**. Data agent **142** may be a software program (e.g., in the form of a set of executable binary files) that executes on the same client computing device **102** as the associated application **110** that data agent **142** is configured to protect. Data agent **142** is generally responsible for managing, initiating, or otherwise assisting in the performance of information management operations in reference to its associated application(s) **110** and corresponding primary data **112** which is generated/accessed by the particular application(s) **110**. For instance, data agent **142** may take part in copying, archiving, migrating, and/or replicating of certain primary data **112** stored in the primary storage device(s) **104**. Data agent **142** may receive control information from storage manager **140**, such as commands to transfer copies of data objects and/or metadata to one or more media agents **144**. Data agent **142** also may compress, deduplicate, and encrypt certain primary data **112**, as well as capture application-related metadata before transmitting the processed data to media agent **144**. Data agent **142** also may receive instructions from storage manager **140** to restore (or assist in restoring) a secondary copy **116** from secondary storage device **108** to primary storage **104**, such that the restored data may be properly accessed by application **110** in a suitable format as though it were primary data **112**.

(84) Each data agent **142** may be specialized for a particular application **110**. For instance, different individual data agents **142** may be designed to handle Microsoft Exchange data, Lotus Notes data, Microsoft Windows file system data, Microsoft Active Directory Objects data, SQL Server data, SharePoint data, Oracle database data, SAP database data, virtual machines and/or associated data, and other types of data. A file system data agent, for example, may handle data files and/or other file system information. If a client computing device **102** has two or more types of data **112**, a specialized data agent **142** may be used for each data type. For example, to backup, migrate, and/or restore all of the data on a Microsoft Exchange server, the client computing device **102** may use: (1) a Microsoft Exchange Mailbox data agent **142** to back up the Exchange mailboxes; (2) a Microsoft Exchange Database data agent **142** to back up the Exchange databases; (3) a Microsoft Exchange Public Folder data agent **142** to back up the Exchange Public Folders; and (4) a Microsoft Windows File System data agent **142** to back up the file system of client computing device **102**. In this example, these specialized data agents **142** are treated as four separate data agents **142** even though they operate on the same client computing device **102**. Other examples may include archive management data agents such as a migration archiver or a compliance archiver, Quick Recovery® agents, and continuous data replication agents. Application-specific data agents **142** can provide improved performance as compared to generic agents. For instance, because application-specific data agents **142** may only handle data for a single software application, the design, operation, and performance of the data agent **142** can be streamlined. The data agent **142** may therefore execute faster and consume less persistent storage and/or operating memory than data agents designed to generically accommodate multiple different software applications **110**.

(85) Each data agent **142** may be configured to access data and/or metadata stored in the primary storage device(s) **104** associated with data agent **142** and its host client computing device **102**, and process the data appropriately. For example, during a secondary copy operation, data agent **142** may arrange or assemble the data and metadata into one or more files having a certain format (e.g., a particular backup or archive format) before transferring the file(s) to a media agent **144** or other component. The file(s) may include a list of files or other metadata. In some embodiments, a data agent **142** may be distributed between client computing device **102** and storage manager **140** (and any other intermediate components) or may be deployed from a remote location or its functions approximated by a remote process that performs some or all of the functions of data agent **142**. In addition, a data agent **142** may perform some functions provided by media agent **144**. Other embodiments may employ one or more generic data agents **142** that can handle and process data from two or more different applications **110**, or that can handle and process multiple data types, instead of or

in addition to using specialized data agents **142**. For example, one generic data agent **142** may be used to back up, migrate and restore Microsoft Exchange Mailbox data and Microsoft Exchange Database data, while another generic data agent may handle Microsoft Exchange Public Folder data and Microsoft Windows File System data.

(86) Media Agents

(87) As noted, off-loading certain responsibilities from client computing devices **102** to intermediate components such as secondary storage computing device(s) **106** and corresponding media agent(s) **144** can provide a number of benefits including improved performance of client computing device **102**, faster and more reliable information management operations, and enhanced scalability. In one example which will be discussed further below, media agent **144** can act as a local cache of recently-copied data and/or metadata stored to secondary storage device(s) **108**, thus improving restore capabilities and performance for the cached data.

(88) Media agent **144** is a component of system **100** and is generally directed by storage manager **140** in creating and restoring secondary copies **116**. Whereas storage manager **140** generally manages system **100** as a whole, media agent **144** provides a portal to certain secondary storage devices **108**, such as by having specialized features for communicating with and accessing certain associated secondary storage device **108**. Media agent **144** may be a software program (e.g., in the form of a set of executable binary files) that executes on a secondary storage computing device **106**. Media agent **144** generally manages, coordinates, and facilitates the transmission of data between a data agent **142** (executing on client computing device **102**) and secondary storage device(s) **108** associated with media agent **144**. For instance, other components in the system may interact with media agent **144** to gain access to data stored on associated secondary storage device(s) **108**, (e.g., to browse, read, write, modify, delete, or restore data). Moreover, media agents **144** can generate and store information relating to characteristics of the stored data and/or metadata, or can generate and store other types of information that generally provides insight into the contents of the secondary storage devices **108**—generally referred to as indexing of the stored secondary copies **116**. Each media agent **144** may operate on a dedicated secondary storage computing device **106**, while in other embodiments a plurality of media agents **144** may operate on the same secondary storage computing device **106**.

(89) A media agent **144** may be associated with a particular secondary storage device **108** if that media agent **144** is capable of one or more of: routing and/or storing data to the particular secondary storage device **108**; coordinating the routing and/or storing of data to the particular secondary storage device **108**; retrieving data from the particular secondary storage device **108**; coordinating the retrieval of data from the particular secondary storage device **108**; and modifying and/or deleting data retrieved from the particular secondary storage device **108**. Media agent **144** in certain embodiments is physically separate from the associated secondary storage device **108**. For instance, a media agent **144** may operate on a secondary storage computing device **106** in a distinct housing, package, and/or location from the associated secondary storage device **108**. In one example, a media agent **144** operates on a first server computer and is in communication with a secondary storage device(s) **108** operating in a separate rack-mounted RAID-based system.

(90) A media agent **144** associated with a particular secondary storage device **108** may instruct secondary storage device **108** to perform an information management task. For instance, a media agent **144** may instruct a tape library to use a robotic arm or other retrieval means to load or eject a certain storage media, and to subsequently archive, migrate, or retrieve data to or from that media, e.g., for the purpose of restoring data to a client computing device **102**. As another example, a secondary storage device **108** may include an array of hard disk drives or solid state drives organized in a RAID configuration, and media agent **144** may forward a logical unit number (LUN) and other appropriate information to the array, which uses the received information to execute the desired secondary copy operation. Media agent **144** may communicate with a secondary storage device **108** via a suitable communications link, such as a SCSI or Fibre Channel link.

(91) Each media agent **144** may maintain an associated media agent database **152**. Media agent database **152** may be stored to a disk or other storage device (not shown) that is local to the secondary storage computing device **106** on which media agent **144** executes. In other cases, media agent

database **152** is stored separately from the host secondary storage computing device **106**. Media agent database **152** can include, among other things, a media agent index **153** (see, e.g., FIG. **1C**). In some cases, media agent index **153** does not form a part of and is instead separate from media agent database **152**.

(92) Media agent index **153** (or “index **153**”) may be a data structure associated with the particular media agent **144** that includes information about the stored data associated with the particular media agent and which may be generated in the course of performing a secondary copy operation or a restore. Index **153** provides a fast and efficient mechanism for locating/browsing secondary copies **116** or other data stored in secondary storage devices **108** without having to access secondary storage device **108** to retrieve the information from there. For instance, for each secondary copy **116**, index **153** may include metadata such as a list of the data objects (e.g., files/subdirectories, database objects, mailbox objects, etc.), a logical path to the secondary copy **116** on the corresponding secondary storage device **108**, location information (e.g., offsets) indicating where the data objects are stored in the secondary storage device **108**, when the data objects were created or modified, etc. Thus, index **153** includes metadata associated with the secondary copies **116** that is readily available for use from media agent **144**. In some embodiments, some or all of the information in index **153** may instead or additionally be stored along with secondary copies **116** in secondary storage device **108**. In some embodiments, a secondary storage device **108** can include sufficient information to enable a “bare metal restore,” where the operating system and/or software applications of a failed client computing device **102** or another target may be automatically restored without manually reinstalling individual software packages (including operating systems).

(93) Because index **153** may operate as a cache, it can also be referred to as an “index cache.” In such cases, information stored in index cache **153** typically comprises data that reflects certain particulars about relatively recent secondary copy operations. After some triggering event, such as after some time elapses or index cache **153** reaches a particular size, certain portions of index cache **153** may be copied or migrated to secondary storage device **108**, e.g., on a least-recently-used basis. This information may be retrieved and uploaded back into index cache **153** or otherwise restored to media agent **144** to facilitate retrieval of data from the secondary storage device(s) **108**. In some embodiments, the cached information may include format or containerization information related to archives or other files stored on storage device(s) **108**.

(94) In some alternative embodiments media agent **144** generally acts as a coordinator or facilitator of secondary copy operations between client computing devices **102** and secondary storage devices **108**, but does not actually write the data to secondary storage device **108**. For instance, storage manager **140** (or media agent **144**) may instruct a client computing device **102** and secondary storage device **108** to communicate with one another directly. In such a case, client computing device **102** transmits data directly or via one or more intermediary components to secondary storage device **108** according to the received instructions, and vice versa. Media agent **144** may still receive, process, and/or maintain metadata related to the secondary copy operations, i.e., may continue to build and maintain index **153**. In these embodiments, payload data can flow through media agent **144** for the purposes of populating index **153**, but not for writing to secondary storage device **108**. Media agent **144** and/or other components such as storage manager **140** may in some cases incorporate additional functionality, such as data classification, content indexing, deduplication, encryption, compression, and the like. Further details regarding these and other functions are described below.

(95) Distributed, Scalable Architecture

(96) As described, certain functions of system **100** can be distributed amongst various physical and/or logical components. For instance, one or more of storage manager **140**, data agents **142**, and media agents **144** may operate on computing devices that are physically separate from one another. This architecture can provide a number of benefits. For instance, hardware and software design choices for each distributed component can be targeted to suit its particular function. The secondary computing devices **106** on which media agents **144** operate can be tailored for interaction with associated secondary storage devices **108** and provide fast index cache operation, among other specific tasks. Similarly, client computing device(s) **102** can be selected to effectively service applications **110** in

order to efficiently produce and store primary data **112**.

(97) Moreover, in some cases, one or more of the individual components of information management system **100** can be distributed to multiple separate computing devices. As one example, for large file systems where the amount of data stored in management database **146** is relatively large, database **146** may be migrated to or may otherwise reside on a specialized database server (e.g., an SQL server) separate from a server that implements the other functions of storage manager **140**. This distributed configuration can provide added protection because database **146** can be protected with standard database utilities (e.g., SQL log shipping or database replication) independent from other functions of storage manager **140**. Database **146** can be efficiently replicated to a remote site for use in the event of a disaster or other data loss at the primary site. Or database **146** can be replicated to another computing device within the same site, such as to a higher performance machine in the event that a storage manager host computing device can no longer service the needs of a growing system **100**.

(98) The distributed architecture also provides scalability and efficient component utilization. FIG. 1D shows an embodiment of information management system **100** including a plurality of client computing devices **102** and associated data agents **142** as well as a plurality of secondary storage computing devices **106** and associated media agents **144**. Additional components can be added or subtracted based on the evolving needs of system **100**. For instance, depending on where bottlenecks are identified, administrators can add additional client computing devices **102**, secondary storage computing devices **106**, and/or secondary storage devices **108**. Moreover, where multiple fungible components are available, load balancing can be implemented to dynamically address identified bottlenecks. As an example, storage manager **140** may dynamically select which media agents **144** and/or secondary storage devices **108** to use for storage operations based on a processing load analysis of media agents **144** and/or secondary storage devices **108**, respectively.

(99) Where system **100** includes multiple media agents **144** (see, e.g., FIG. 1D), a first media agent **144** may provide failover functionality for a second failed media agent **144**. In addition, media agents **144** can be dynamically selected to provide load balancing. Each client computing device **102** can communicate with, among other components, any of the media agents **144**, e.g., as directed by storage manager **140**. And each media agent **144** may communicate with, among other components, any of secondary storage devices **108**, e.g., as directed by storage manager **140**. Thus, operations can be routed to secondary storage devices **108** in a dynamic and highly flexible manner, to provide load balancing, failover, etc. Further examples of scalable systems capable of dynamic storage operations, load balancing, and failover are provided in U.S. Pat. No. 7,246,207.

(100) While distributing functionality amongst multiple computing devices can have certain advantages, in other contexts it can be beneficial to consolidate functionality on the same computing device. In alternative configurations, certain components may reside and execute on the same computing device. As such, in other embodiments, one or more of the components shown in FIG. 1C may be implemented on the same computing device. In one configuration, a storage manager **140**, one or more data agents **142**, and/or one or more media agents **144** are all implemented on the same computing device. In other embodiments, one or more data agents **142** and one or more media agents **144** are implemented on the same computing device, while storage manager **140** is implemented on a separate computing device, etc. without limitation.

(101) Exemplary Types of Information Management Operations, Including Storage Operations

(102) In order to protect and leverage stored data, system **100** can be configured to perform a variety of information management operations, which may also be referred to in some cases as storage management operations or storage operations. These operations can generally include (i) data movement operations, (ii) processing and data manipulation operations, and (iii) analysis, reporting, and management operations.

(103) Data Movement Operations, Including Secondary Copy Operations

(104) Data movement operations are generally storage operations that involve the copying or migration of data between different locations in system **100**. For example, data movement operations can include operations in which stored data is copied, migrated, or otherwise transferred from one or more first storage devices to one or more second storage devices, such as from primary storage

device(s) **104** to secondary storage device(s) **108**, from secondary storage device(s) **108** to different secondary storage device(s) **108**, from secondary storage devices **108** to primary storage devices **104**, or from primary storage device(s) **104** to different primary storage device(s) **104**, or in some cases within the same primary storage device **104** such as within a storage array.

(105) Data movement operations can include by way of example, backup operations, archive operations, information lifecycle management operations such as hierarchical storage management operations, replication operations (e.g., continuous data replication), snapshot operations, deduplication or single-instancing operations, auxiliary copy operations, disaster-recovery copy operations, and the like. As will be discussed, some of these operations do not necessarily create distinct copies. Nonetheless, some or all of these operations are generally referred to as “secondary copy operations” for simplicity, because they involve secondary copies. Data movement also comprises restoring secondary copies.

(106) Backup Operations

(107) A backup operation creates a copy of a version of primary data **112** at a particular point in time (e.g., one or more files or other data units). Each subsequent backup copy **116** (which is a form of secondary copy **116**) may be maintained independently of the first. A backup generally involves maintaining a version of the copied primary data **112** as well as backup copies **116**. Further, a backup copy in some embodiments is generally stored in a form that is different from the native format, e.g., a backup format. This contrasts to the version in primary data **112** which may instead be stored in a format native to the source application(s) **110**. In various cases, backup copies can be stored in a format in which the data is compressed, encrypted, deduplicated, and/or otherwise modified from the original native application format. For example, a backup copy may be stored in a compressed backup format that facilitates efficient long-term storage. Backup copies **116** can have relatively long retention periods as compared to primary data **112**, which is generally highly changeable. Backup copies **116** may be stored on media with slower retrieval times than primary storage device **104**. Some backup copies may have shorter retention periods than some other types of secondary copies **116**, such as archive copies (described below). Backups may be stored at an offsite location.

(108) Backup operations can include full backups, differential backups, incremental backups, “synthetic full” backups, and/or creating a “reference copy.” A full backup (or “standard full backup”) in some embodiments is generally a complete image of the data to be protected. However, because full backup copies can consume a relatively large amount of storage, it can be useful to use a full backup copy as a baseline and only store changes relative to the full backup copy afterwards.

(109) A differential backup operation (or cumulative incremental backup operation) tracks and stores changes that occurred since the last full backup. Differential backups can grow quickly in size, but can restore relatively efficiently because a restore can be completed in some cases using only the full backup copy and the latest differential copy.

(110) An incremental backup operation generally tracks and stores changes since the most recent backup copy of any type, which can greatly reduce storage utilization. In some cases, however, restoring can be lengthy compared to full or differential backups because completing a restore operation may involve accessing a full backup in addition to multiple incremental backups.

(111) Synthetic full backups generally consolidate data without directly backing up data from the client computing device. A synthetic full backup is created from the most recent full backup (i.e., standard or synthetic) and subsequent incremental and/or differential backups. The resulting synthetic full backup is identical to what would have been created had the last backup for the subclient been a standard full backup. Unlike standard full, incremental, and differential backups, however, a synthetic full backup does not actually transfer data from primary storage to the backup media, because it operates as a backup consolidator. A synthetic full backup extracts the index data of each participating subclient. Using this index data and the previously backed up user data images, it builds new full backup images (e.g., bitmaps), one for each subclient. The new backup images consolidate the index and user data stored in the related incremental, differential, and previous full backups into a synthetic backup file that fully represents the subclient (e.g., via pointers) but does not comprise all its constituent data.

(112) Any of the above types of backup operations can be at the volume level, file level, or block level.

Volume level backup operations generally involve copying of a data volume (e.g., a logical disk or partition) as a whole. In a file-level backup, information management system **100** generally tracks changes to individual files and includes copies of files in the backup copy. For block-level backups, files are broken into constituent blocks, and changes are tracked at the block level. Upon restore, system **100** reassembles the blocks into files in a transparent fashion. Far less data may actually be transferred and copied to secondary storage devices **108** during a file-level copy than a volume-level copy. Likewise, a block-level copy may transfer less data than a file-level copy, resulting in faster execution. However, restoring a relatively higher-granularity copy can result in longer restore times. For instance, when restoring a block-level copy, the process of locating and retrieving constituent blocks can sometimes take longer than restoring file-level backups.

(113) A reference copy may comprise copy(ies) of selected objects from backed up data, typically to help organize data by keeping contextual information from multiple sources together, and/or help retain specific data for a longer period of time, such as for legal hold needs. A reference copy generally maintains data integrity, and when the data is restored, it may be viewed in the same format as the source data. In some embodiments, a reference copy is based on a specialized client, individual subclient and associated information management policies (e.g., storage policy, retention policy, etc.) that are administered within system **100**.

(114) Archive Operations

(115) Because backup operations generally involve maintaining a version of the copied primary data **112** and also maintaining backup copies in secondary storage device(s) **108**, they can consume significant storage capacity. To reduce storage consumption, an archive operation according to certain embodiments creates an archive copy **116** by both copying and removing source data. Or, seen another way, archive operations can involve moving some or all of the source data to the archive destination. Thus, data satisfying criteria for removal (e.g., data of a threshold age or size) may be removed from source storage. The source data may be primary data **112** or a secondary copy **116**, depending on the situation. As with backup copies, archive copies can be stored in a format in which the data is compressed, encrypted, deduplicated, and/or otherwise modified from the format of the original application or source copy. In addition, archive copies may be retained for relatively long periods of time (e.g., years) and, in some cases are never deleted. In certain embodiments, archive copies may be made and kept for extended periods in order to meet compliance regulations.

(116) Archiving can also serve the purpose of freeing up space in primary storage device(s) **104** and easing the demand on computational resources on client computing device **102**. Similarly, when a secondary copy **116** is archived, the archive copy can therefore serve the purpose of freeing up space in the source secondary storage device(s) **108**. Examples of data archiving operations are provided in U.S. Pat. No. 7,107,298.

(117) Snapshot Operations

(118) Snapshot operations can provide a relatively lightweight, efficient mechanism for protecting data. From an end-user viewpoint, a snapshot may be thought of as an “instant” image of primary data **112** at a given point in time, and may include state and/or status information relative to an application **110** that creates/manages primary data **112**. In one embodiment, a snapshot may generally capture the directory structure of an object in primary data **112** such as a file or volume or other data set at a particular moment in time and may also preserve file attributes and contents. A snapshot in some cases is created relatively quickly, e.g., substantially instantly, using a minimum amount of file space, but may still function as a conventional file system backup.

(119) A “hardware snapshot” (or “hardware-based snapshot”) operation occurs where a target storage device (e.g., a primary storage device **104** or a secondary storage device **108**) performs the snapshot operation in a self-contained fashion, substantially independently, using hardware, firmware and/or software operating on the storage device itself. For instance, the storage device may perform snapshot operations generally without intervention or oversight from any of the other components of the system **100**, e.g., a storage array may generate an “array-created” hardware snapshot and may also manage its storage, integrity, versioning, etc. In this manner, hardware snapshots can off-load other components of system **100** from snapshot processing. An array may receive a request from another component to take

a snapshot and then proceed to execute the “hardware snapshot” operations autonomously, preferably reporting success to the requesting component.

(120) A “software snapshot” (or “software-based snapshot”) operation, on the other hand, occurs where a component in system **100** (e.g., client computing device **102**, etc.) implements a software layer that manages the snapshot operation via interaction with the target storage device. For instance, the component executing the snapshot management software layer may derive a set of pointers and/or data that represents the snapshot. The snapshot management software layer may then transmit the same to the target storage device, along with appropriate instructions for writing the snapshot. One example of a software snapshot product is Microsoft Volume Snapshot Service (VSS), which is part of the Microsoft Windows operating system.

(121) Some types of snapshots do not actually create another physical copy of all the data as it existed at the particular point in time, but may simply create pointers that map files and directories to specific memory locations (e.g., to specific disk blocks) where the data resides as it existed at the particular point in time. For example, a snapshot copy may include a set of pointers derived from the file system or from an application. In some other cases, the snapshot may be created at the block-level, such that creation of the snapshot occurs without awareness of the file system. Each pointer points to a respective stored data block, so that collectively, the set of pointers reflect the storage location and state of the data object (e.g., file(s) or volume(s) or data set(s)) at the point in time when the snapshot copy was created.

(122) An initial snapshot may use only a small amount of disk space needed to record a mapping or other data structure representing or otherwise tracking the blocks that correspond to the current state of the file system. Additional disk space is usually required only when files and directories change later on. Furthermore, when files change, typically only the pointers which map to blocks are copied, not the blocks themselves. For example for “copy-on-write” snapshots, when a block changes in primary storage, the block is copied to secondary storage or cached in primary storage before the block is overwritten in primary storage, and the pointer to that block is changed to reflect the new location of that block. The snapshot mapping of file system data may also be updated to reflect the changed block(s) at that particular point in time. In some other cases, a snapshot includes a full physical copy of all or substantially all of the data represented by the snapshot. Further examples of snapshot operations are provided in U.S. Pat. No. 7,529,782. A snapshot copy in many cases can be made quickly and without significantly impacting primary computing resources because large amounts of data need not be copied or moved. In some embodiments, a snapshot may exist as a virtual file system, parallel to the actual file system. Users in some cases gain read-only access to the record of files and directories of the snapshot. By electing to restore primary data **112** from a snapshot taken at a given point in time, users may also return the current file system to the state of the file system that existed when the snapshot was taken.

(123) Replication Operations

(124) Replication is another type of secondary copy operation. Some types of secondary copies **116** periodically capture images of primary data **112** at particular points in time (e.g., backups, archives, and snapshots). However, it can also be useful for recovery purposes to protect primary data **112** in a more continuous fashion, by replicating primary data **112** substantially as changes occur. In some cases a replication copy can be a mirror copy, for instance, where changes made to primary data **112** are mirrored or substantially immediately copied to another location (e.g., to secondary storage device(s) **108**). By copying each write operation to the replication copy, two storage systems are kept synchronized or substantially synchronized so that they are virtually identical at approximately the same time. Where entire disk volumes are mirrored, however, mirroring can require significant amount of storage space and utilizes a large amount of processing resources.

(125) According to some embodiments, secondary copy operations are performed on replicated data that represents a recoverable state, or “known good state” of a particular application running on the source system. For instance, in certain embodiments, known good replication copies may be viewed as copies of primary data **112**. This feature allows the system to directly access, copy, restore, back up, or otherwise manipulate the replication copies as if they were the “live” primary data **112**. This can

reduce access time, storage utilization, and impact on source applications **110**, among other benefits. Based on known good state information, system **100** can replicate sections of application data that represent a recoverable state rather than rote copying of blocks of data. Examples of replication operations (e.g., continuous data replication) are provided in U.S. Pat. No. 7,617,262.

(126) Deduplication/Single-Instancing Operations

(127) Deduplication or single-instance storage is useful to reduce the amount of non-primary data. For instance, some or all of the above-described secondary copy operations can involve deduplication in some fashion. New data is read, broken down into data portions of a selected granularity (e.g., sub-file level blocks, files, etc.), compared with corresponding portions that are already in secondary storage, and only new/changed portions are stored. Portions that already exist are represented as pointers to the already-stored data. Thus, a deduplicated secondary copy **116** may comprise actual data portions copied from primary data **112** and may further comprise pointers to already-stored data, which is generally more storage-efficient than a full copy.

(128) In order to streamline the comparison process, system **100** may calculate and/or store signatures (e.g., hashes or cryptographically unique IDs) corresponding to the individual source data portions and compare the signatures to already-stored data signatures, instead of comparing entire data portions. In some cases, only a single instance of each data portion is stored, and deduplication operations may therefore be referred to interchangeably as “single-instancing” operations. Depending on the implementation, however, deduplication operations can store more than one instance of certain data portions, yet still significantly reduce stored-data redundancy. Depending on the embodiment, deduplication portions such as data blocks can be of fixed or variable length. Using variable length blocks can enhance deduplication by responding to changes in the data stream, but can involve more complex processing. In some cases, system **100** utilizes a technique for dynamically aligning deduplication blocks based on changing content in the data stream, as described in U.S. Pat. No. 8,364,652.

(129) System **100** can deduplicate in a variety of manners at a variety of locations. For instance, in some embodiments, system **100** implements “target-side” deduplication by deduplicating data at the media agent **144** after being received from data agent **142**. In some such cases, media agents **144** are generally configured to manage the deduplication process. For instance, one or more of the media agents **144** maintain a corresponding deduplication database that stores deduplication information (e.g., data block signatures). Examples of such a configuration are provided in U.S. Pat. No. 9,020,900. Instead of or in combination with “target-side” deduplication, “source-side” (or “client-side”) deduplication can also be performed, e.g., to reduce the amount of data to be transmitted by data agent **142** to media agent **144**. Storage manager **140** may communicate with other components within system **100** via network protocols and cloud service provider APIs to facilitate cloud-based deduplication/single instancing, as exemplified in U.S. Pat. No. 8,954,446. Some other deduplication/single instancing techniques are described in U.S. Pat. Pub. No. 2006/0224846 and in U.S. Pat. No. 9,098,495.

(130) Information Lifecycle Management and Hierarchical Storage Management

(131) In some embodiments, files and other data over their lifetime move from more expensive quick-access storage to less expensive slower-access storage. Operations associated with moving data through various tiers of storage are sometimes referred to as information lifecycle management (ILM) operations.

(132) One type of ILM operation is a hierarchical storage management (HSM) operation, which generally automatically moves data between classes of storage devices, such as from high-cost to low-cost storage devices. For instance, an HSM operation may involve movement of data from primary storage devices **104** to secondary storage devices **108**, or between tiers of secondary storage devices **108**. With each tier, the storage devices may be progressively cheaper, have relatively slower access/restore times, etc. For example, movement of data between tiers may occur as data becomes less important over time. In some embodiments, an HSM operation is similar to archiving in that creating an HSM copy may (though not always) involve deleting some of the source data, e.g., according to one or more criteria related to the source data. For example, an HSM copy may include

primary data **112** or a secondary copy **116** that exceeds a given size threshold or a given age threshold. Often, and unlike some types of archive copies, HSM data that is removed or aged from the source is replaced by a logical reference pointer or stub. The reference pointer or stub can be stored in the primary storage device **104** or other source storage device, such as a secondary storage device **108** to replace the deleted source data and to point to or otherwise indicate the new location in (another) secondary storage device **108**.

(133) For example, files are generally moved between higher and lower cost storage depending on how often the files are accessed. When a user requests access to HSM data that has been removed or migrated, system **100** uses the stub to locate the data and can make recovery of the data appear transparent, even though the HSM data may be stored at a location different from other source data. In this manner, the data appears to the user (e.g., in file system browsing windows and the like) as if it still resides in the source location (e.g., in a primary storage device **104**). The stub may include metadata associated with the corresponding data, so that a file system and/or application can provide some information about the data object and/or a limited-functionality version (e.g., a preview) of the data object.

(134) An HSM copy may be stored in a format other than the native application format (e.g., compressed, encrypted, deduplicated, and/or otherwise modified). In some cases, copies which involve the removal of data from source storage and the maintenance of stub or other logical reference information on source storage may be referred to generally as “online archive copies.” On the other hand, copies which involve the removal of data from source storage without the maintenance of stub or other logical reference information on source storage may be referred to as “off-line archive copies.” Examples of HSM and ILM techniques are provided in U.S. Pat. No. 7,343,453.

(135) Auxiliary Copy Operations

(136) An auxiliary copy is generally a copy of an existing secondary copy **116**. For instance, an initial secondary copy **116** may be derived from primary data **112** or from data residing in secondary storage subsystem **118**, whereas an auxiliary copy is generated from the initial secondary copy **116**. Auxiliary copies provide additional standby copies of data and may reside on different secondary storage devices **108** than the initial secondary copies **116**. Thus, auxiliary copies can be used for recovery purposes if initial secondary copies **116** become unavailable. Exemplary auxiliary copy techniques are described in further detail in U.S. Pat. No. 8,230,195.

(137) Disaster-Recovery Copy Operations

(138) System **100** may also make and retain disaster recovery copies, often as secondary, high-availability disk copies. System **100** may create secondary copies and store them at disaster recovery locations using auxiliary copy or replication operations, such as continuous data replication technologies. Depending on the particular data protection goals, disaster recovery locations can be remote from the client computing devices **102** and primary storage devices **104**, remote from some or all of the secondary storage devices **108**, or both.

(139) Data Manipulation, Including Encryption and Compression

(140) Data manipulation and processing may include encryption and compression as well as integrity marking and checking, formatting for transmission, formatting for storage, etc. Data may be manipulated “client-side” by data agent **142** as well as “target-side” by media agent **144** in the course of creating secondary copy **116**, or conversely in the course of restoring data from secondary to primary.

(141) Encryption Operations

(142) System **100** in some cases is configured to process data (e.g., files or other data objects, primary data **112**, secondary copies **116**, etc.), according to an appropriate encryption algorithm (e.g., Blowfish, Advanced Encryption Standard (AES), Triple Data Encryption Standard (3-DES), etc.) to limit access and provide data security. System **100** in some cases encrypts the data at the client level, such that client computing devices **102** (e.g., data agents **142**) encrypt the data prior to transferring it to other components, e.g., before sending the data to media agents **144** during a secondary copy operation. In such cases, client computing device **102** may maintain or have access to an encryption key or passphrase for decrypting the data upon restore. Encryption can also occur when media agent **144**

creates auxiliary copies or archive copies. Encryption may be applied in creating a secondary copy **116** of a previously unencrypted secondary copy **116**, without limitation. In further embodiments, secondary storage devices **108** can implement built-in, high performance hardware-based encryption.

(143) Compression Operations

(144) Similar to encryption, system **100** may also or alternatively compress data in the course of generating a secondary copy **116**. Compression encodes information such that fewer bits are needed to represent the information as compared to the original representation. Compression techniques are well known in the art. Compression operations may apply one or more data compression algorithms. Compression may be applied in creating a secondary copy **116** of a previously uncompressed secondary copy, e.g., when making archive copies or disaster recovery copies. The use of compression may result in metadata that specifies the nature of the compression, so that data may be uncompressed on restore if appropriate.

(145) Data Analysis, Reporting, and Management Operations

(146) Data analysis, reporting, and management operations can differ from data movement operations in that they do not necessarily involve copying, migration or other transfer of data between different locations in the system. For instance, data analysis operations may involve processing (e.g., offline processing) or modification of already stored primary data **112** and/or secondary copies **116**. However, in some embodiments data analysis operations are performed in conjunction with data movement operations. Some data analysis operations include content indexing operations and classification operations which can be useful in leveraging data under management to enhance search and other features.

(147) Classification Operations/Content Indexing

(148) In some embodiments, information management system **100** analyzes and indexes characteristics, content, and metadata associated with primary data **112** (“online content indexing”) and/or secondary copies **116** (“off-line content indexing”). Content indexing can identify files or other data objects based on content (e.g., user-defined keywords or phrases, other keywords/phrases that are not defined by a user, etc.), and/or metadata (e.g., email metadata such as “to,” “from,” “cc,” “bcc,” attachment name, received time, etc.). Content indexes may be searched and search results may be restored.

(149) System **100** generally organizes and catalogues the results into a content index, which may be stored within media agent database **152**, for example. The content index can also include the storage locations of or pointer references to indexed data in primary data **112** and/or secondary copies **116**. Results may also be stored elsewhere in system **100** (e.g., in primary storage device **104** or in secondary storage device **108**). Such content index data provides storage manager **140** or other components with an efficient mechanism for locating primary data **112** and/or secondary copies **116** of data objects that match particular criteria, thus greatly increasing the search speed capability of system **100**. For instance, search criteria can be specified by a user through user interface **158** of storage manager **140**. Moreover, when system **100** analyzes data and/or metadata in secondary copies **116** to create an “off-line content index,” this operation has no significant impact on the performance of client computing devices **102** and thus does not take a toll on the production environment. Examples of content indexing techniques are provided in U.S. Pat. No. 8,170,995.

(150) One or more components, such as a content index engine, can be configured to scan data and/or associated metadata for classification purposes to populate a database (or other data structure) of information, which can be referred to as a “data classification database” or a “metabase.” Depending on the embodiment, the data classification database(s) can be organized in a variety of different ways, including centralization, logical sub-divisions, and/or physical sub-divisions. For instance, one or more data classification databases may be associated with different subsystems or tiers within system **100**. As an example, there may be a first metabase associated with primary storage subsystem **117** and a second metabase associated with secondary storage subsystem **118**. In other cases, metabase(s) may be associated with individual components, e.g., client computing devices **102** and/or media agents **144**. In some embodiments, a data classification database may reside as one or more data structures within management database **146**, may be otherwise associated with storage manager **140**, and/or may reside

as a separate component. In some cases, metabase(s) may be included in separate database(s) and/or on separate storage device(s) from primary data **112** and/or secondary copies **116**, such that operations related to the metabase(s) do not significantly impact performance on other components of system **100**. In other cases, metabase(s) may be stored along with primary data **112** and/or secondary copies **116**. Files or other data objects can be associated with identifiers (e.g., tag entries, etc.) to facilitate searches of stored data objects. Among a number of other benefits, the metabase can also allow efficient, automatic identification of files or other data objects to associate with secondary copy or other information management operations. For instance, a metabase can dramatically improve the speed with which system **100** can search through and identify data as compared to other approaches that involve scanning an entire file system. Examples of metabases and data classification operations are provided in U.S. Pat. Nos. 7,734,669 and 7,747,579.

(151) Management and Reporting Operations

(152) Certain embodiments leverage the integrated ubiquitous nature of system **100** to provide useful system-wide management and reporting. Operations management can generally include monitoring and managing the health and performance of system **100** by, without limitation, performing error tracking, generating granular storage/performance metrics (e.g., job success/failure information, deduplication efficiency, etc.), generating storage modeling and costing information, and the like. As an example, storage manager **140** or another component in system **100** may analyze traffic patterns and suggest and/or automatically route data to minimize congestion. In some embodiments, the system can generate predictions relating to storage operations or storage operation information. Such predictions, which may be based on a trending analysis, may predict various network operations or resource usage, such as network traffic levels, storage media use, use of bandwidth of communication links, use of media agent components, etc. Further examples of traffic analysis, trend analysis, prediction generation, and the like are described in U.S. Pat. No. 7,343,453.

(153) In some configurations having a hierarchy of storage operation cells, a master storage manager **140** may track the status of subordinate cells, such as the status of jobs, system components, system resources, and other items, by communicating with storage managers **140** (or other components) in the respective storage operation cells. Moreover, the master storage manager **140** may also track status by receiving periodic status updates from the storage managers **140** (or other components) in the respective cells regarding jobs, system components, system resources, and other items. In some embodiments, a master storage manager **140** may store status information and other information regarding its associated storage operation cells and other system information in its management database **146** and/or index **150** (or in another location). The master storage manager **140** or other component may also determine whether certain storage-related or other criteria are satisfied, and may perform an action or trigger event (e.g., data migration) in response to the criteria being satisfied, such as where a storage threshold is met for a particular volume, or where inadequate protection exists for certain data. For instance, data from one or more storage operation cells is used to dynamically and automatically mitigate recognized risks, and/or to advise users of risks or suggest actions to mitigate these risks. For example, an information management policy may specify certain requirements (e.g., that a storage device should maintain a certain amount of free space, that secondary copies should occur at a particular interval, that data should be aged and migrated to other storage after a particular period, that data on a secondary volume should always have a certain level of availability and be restorable within a given time period, that data on a secondary volume may be mirrored or otherwise migrated to a specified number of other volumes, etc.). If a risk condition or other criterion is triggered, the system may notify the user of these conditions and may suggest (or automatically implement) a mitigation action to address the risk. For example, the system may indicate that data from a primary copy **112** should be migrated to a secondary storage device **108** to free up space on primary storage device **104**. Examples of the use of risk factors and other triggering criteria are described in U.S. Pat. No. 7,343,453.

(154) In some embodiments, system **100** may also determine whether a metric or other indication satisfies particular storage criteria sufficient to perform an action. For example, a storage policy or other definition might indicate that a storage manager **140** should initiate a particular action if a

storage metric or other indication drops below or otherwise fails to satisfy specified criteria such as a threshold of data protection. In some embodiments, risk factors may be quantified into certain measurable service or risk levels. For example, certain applications and associated data may be considered to be more important relative to other data and services. Financial compliance data, for example, may be of greater importance than marketing materials, etc. Network administrators may assign priority values or “weights” to certain data and/or applications corresponding to the relative importance. The level of compliance of secondary copy operations specified for these applications may also be assigned a certain value. Thus, the health, impact, and overall importance of a service may be determined, such as by measuring the compliance value and calculating the product of the priority value and the compliance value to determine the “service level” and comparing it to certain operational thresholds to determine whether it is acceptable. Further examples of the service level determination are provided in U.S. Pat. No. 7,343,453.

(155) System **100** may additionally calculate data costing and data availability associated with information management operation cells. For instance, data received from a cell may be used in conjunction with hardware-related information and other information about system elements to determine the cost of storage and/or the availability of particular data. Exemplary information generated could include how fast a particular department is using up available storage space, how long data would take to recover over a particular pathway from a particular secondary storage device, costs over time, etc. Moreover, in some embodiments, such information may be used to determine or predict the overall cost associated with the storage of certain information. The cost associated with hosting a certain application may be based, at least in part, on the type of media on which the data resides, for example. Storage devices may be assigned to a particular cost categories, for example. Further examples of costing techniques are described in U.S. Pat. No. 7,343,453.

(156) Any of the above types of information (e.g., information related to trending, predictions, job, cell or component status, risk, service level, costing, etc.) can generally be provided to users via user interface **158** in a single integrated view or console (not shown). Report types may include: scheduling, event management, media management and data aging. Available reports may also include backup history, data aging history, auxiliary copy history, job history, library and drive, media in library, restore history, and storage policy, etc., without limitation. Such reports may be specified and created at a certain point in time as a system analysis, forecasting, or provisioning tool. Integrated reports may also be generated that illustrate storage and performance metrics, risks and storage costing information. Moreover, users may create their own reports based on specific needs. User interface **158** can include an option to graphically depict the various components in the system using appropriate icons. As one example, user interface **158** may provide a graphical depiction of primary storage devices **104**, secondary storage devices **108**, data agents **142** and/or media agents **144**, and their relationship to one another in system **100**.

(157) In general, the operations management functionality of system **100** can facilitate planning and decision-making. For example, in some embodiments, a user may view the status of some or all jobs as well as the status of each component of information management system **100**. Users may then plan and make decisions based on this data. For instance, a user may view high-level information regarding secondary copy operations for system **100**, such as job status, component status, resource status (e.g., communication pathways, etc.), and other information. The user may also drill down or use other means to obtain more detailed information regarding a particular component, job, or the like. Further examples are provided in U.S. Pat. No. 7,343,453.

(158) System **100** can also be configured to perform system-wide e-discovery operations in some embodiments. In general, e-discovery operations provide a unified collection and search capability for data in the system, such as data stored in secondary storage devices **108** (e.g., backups, archives, or other secondary copies **116**). For example, system **100** may construct and maintain a virtual repository for data stored in system **100** that is integrated across source applications **110**, different storage device types, etc. According to some embodiments, e-discovery utilizes other techniques described herein, such as data classification and/or content indexing.

(159) Information Management Policies

(160) An information management policy **148** can include a data structure or other information source that specifies a set of parameters (e.g., criteria and rules) associated with secondary copy and/or other information management operations.

(161) One type of information management policy **148** is a “storage policy.” According to certain embodiments, a storage policy generally comprises a data structure or other information source that defines (or includes information sufficient to determine) a set of preferences or other criteria for performing information management operations. Storage policies can include one or more of the following: (1) what data will be associated with the storage policy, e.g., subclient; (2) a destination to which the data will be stored; (3) datapath information specifying how the data will be communicated to the destination; (4) the type of secondary copy operation to be performed; and (5) retention information specifying how long the data will be retained at the destination (see, e.g., FIG. 1E). Data associated with a storage policy can be logically organized into subclients, which may represent primary data **112** and/or secondary copies **116**. A subclient may represent static or dynamic associations of portions of a data volume. Subclients may represent mutually exclusive portions. Thus, in certain embodiments, a portion of data may be given a label and the association is stored as a static entity in an index, database or other storage location. Subclients may also be used as an effective administrative scheme of organizing data according to data type, department within the enterprise, storage preferences, or the like. Depending on the configuration, subclients can correspond to files, folders, virtual machines, databases, etc. In one exemplary scenario, an administrator may find it preferable to separate e-mail data from financial data using two different subclients.

(162) A storage policy can define where data is stored by specifying a target or destination storage device (or group of storage devices). For instance, where the secondary storage device **108** includes a group of disk libraries, the storage policy may specify a particular disk library for storing the subclients associated with the policy. As another example, where the secondary storage devices **108** include one or more tape libraries, the storage policy may specify a particular tape library for storing the subclients associated with the storage policy, and may also specify a drive pool and a tape pool defining a group of tape drives and a group of tapes, respectively, for use in storing the subclient data. While information in the storage policy can be statically assigned in some cases, some or all of the information in the storage policy can also be dynamically determined based on criteria set forth in the storage policy. For instance, based on such criteria, a particular destination storage device(s) or other parameter of the storage policy may be determined based on characteristics associated with the data involved in a particular secondary copy operation, device availability (e.g., availability of a secondary storage device **108** or a media agent **144**), network status and conditions (e.g., identified bottlenecks), user credentials, and the like.

(163) Datapath information can also be included in the storage policy. For instance, the storage policy may specify network pathways and components to utilize when moving the data to the destination storage device(s). In some embodiments, the storage policy specifies one or more media agents **144** for conveying data associated with the storage policy between the source and destination. A storage policy can also specify the type(s) of associated operations, such as backup, archive, snapshot, auxiliary copy, or the like. Furthermore, retention parameters can specify how long the resulting secondary copies **116** will be kept (e.g., a number of days, months, years, etc.), perhaps depending on organizational needs and/or compliance criteria.

(164) When adding a new client computing device **102**, administrators can manually configure information management policies **148** and/or other settings, e.g., via user interface **158**. However, this can be an involved process resulting in delays, and it may be desirable to begin data protection operations quickly, without awaiting human intervention. Thus, in some embodiments, system **100** automatically applies a default configuration to client computing device **102**. As one example, when one or more data agent(s) **142** are installed on a client computing device **102**, the installation script may register the client computing device **102** with storage manager **140**, which in turn applies the default configuration to the new client computing device **102**. In this manner, data protection operations can begin substantially immediately. The default configuration can include a default storage policy, for example, and can specify any appropriate information sufficient to begin data protection

operations. This can include a type of data protection operation, scheduling information, a target secondary storage device **108**, data path information (e.g., a particular media agent **144**), and the like. (165) Another type of information management policy **148** is a “scheduling policy,” which specifies when and how often to perform operations. Scheduling parameters may specify with what frequency (e.g., hourly, weekly, daily, event-based, etc.) or under what triggering conditions secondary copy or other information management operations are to take place. Scheduling policies in some cases are associated with particular components, such as a subclient, client computing device **102**, and the like. (166) Another type of information management policy **148** is an “audit policy” (or “security policy”), which comprises preferences, rules and/or criteria that protect sensitive data in system **100**. For example, an audit policy may define “sensitive objects” which are files or data objects that contain particular keywords (e.g., “confidential,” or “privileged”) and/or are associated with particular keywords (e.g., in metadata) or particular flags (e.g., in metadata identifying a document or email as personal, confidential, etc.). An audit policy may further specify rules for handling sensitive objects. As an example, an audit policy may require that a reviewer approve the transfer of any sensitive objects to a cloud storage site, and that if approval is denied for a particular sensitive object, the sensitive object should be transferred to a local primary storage device **104** instead. To facilitate this approval, the audit policy may further specify how a secondary storage computing device **106** or other system component should notify a reviewer that a sensitive object is slated for transfer. (167) Another type of information management policy **148** is a “provisioning policy,” which can include preferences, priorities, rules, and/or criteria that specify how client computing devices **102** (or groups thereof) may utilize system resources, such as available storage on cloud storage and/or network bandwidth. A provisioning policy specifies, for example, data quotas for particular client computing devices **102** (e.g., a number of gigabytes that can be stored monthly, quarterly or annually). Storage manager **140** or other components may enforce the provisioning policy. For instance, media agents **144** may enforce the policy when transferring data to secondary storage devices **108**. If a client computing device **102** exceeds a quota, a budget for the client computing device **102** (or associated department) may be adjusted accordingly or an alert may trigger. (168) While the above types of information management policies **148** are described as separate policies, one or more of these can be generally combined into a single information management policy **148**. For instance, a storage policy may also include or otherwise be associated with one or more scheduling, audit, or provisioning policies or operational parameters thereof. Moreover, while storage policies are typically associated with moving and storing data, other policies may be associated with other types of information management operations. The following is a non-exhaustive list of items that information management policies **148** may specify: schedules or other timing information, e.g., specifying when and/or how often to perform information management operations; the type of secondary copy **116** and/or copy format (e.g., snapshot, backup, archive, HSM, etc.); a location or a class or quality of storage for storing secondary copies **116** (e.g., one or more particular secondary storage devices **108**); preferences regarding whether and how to encrypt, compress, deduplicate, or otherwise modify or transform secondary copies **116**; which system components and/or network pathways (e.g., preferred media agents **144**) should be used to perform secondary storage operations; resource allocation among different computing devices or other system components used in performing information management operations (e.g., bandwidth allocation, available storage capacity, etc.); whether and how to synchronize or otherwise distribute files or other data objects across multiple computing devices or hosted services; and retention information specifying the length of time primary data **112** and/or secondary copies **116** should be retained, e.g., in a particular class or tier of storage devices, or within the system **100**. (169) Information management policies **148** can additionally specify or depend on historical or current criteria that may be used to determine which rules to apply to a particular data object, system component, or information management operation, such as: frequency with which primary data **112** or a secondary copy **116** of a data object or metadata has been or is predicted to be used, accessed, or modified; time-related factors (e.g., aging information such as time since the creation or modification of a data object); deduplication information (e.g., hashes, data blocks, deduplication block size,

deduplication efficiency or other metrics); an estimated or historic usage or cost associated with different components (e.g., with secondary storage devices **108**); the identity of users, applications **110**, client computing devices **102** and/or other computing devices that created, accessed, modified, or otherwise utilized primary data **112** or secondary copies **116**; a relative sensitivity (e.g., confidentiality, importance) of a data object, e.g., as determined by its content and/or metadata; the current or historical storage capacity of various storage devices; the current or historical network capacity of network pathways connecting various components within the storage operation cell; access control lists or other security information; and the content of a particular data object (e.g., its textual content) or of metadata associated with the data object.

Exemplary Storage Policy and Secondary Copy Operations

(170) FIG. **1E** includes a data flow diagram depicting performance of secondary copy operations by an embodiment of information management system **100**, according to an exemplary storage policy **148A**. System **100** includes a storage manager **140**, a client computing device **102** having a file system data agent **142A** and an email data agent **142B** operating thereon, a primary storage device **104**, two media agents **144A**, **144B**, and two secondary storage devices **108**: a disk library **108A** and a tape library **108B**. As shown, primary storage device **104** includes primary data **112A**, which is associated with a logical grouping of data associated with a file system (“file system subclient”), and primary data **112B**, which is a logical grouping of data associated with email (“email subclient”). The techniques described with respect to FIG. **1E** can be utilized in conjunction with data that is otherwise organized as well.

(171) As indicated by the dashed box, the second media agent **144B** and tape library **108B** are “off-site,” and may be remotely located from the other components in system **100** (e.g., in a different city, office building, etc.). Indeed, “off-site” may refer to a magnetic tape located in remote storage, which must be manually retrieved and loaded into a tape drive to be read. In this manner, information stored on the tape library **108B** may provide protection in the event of a disaster or other failure at the main site(s) where data is stored.

(172) The file system subclient **112A** in certain embodiments generally comprises information generated by the file system and/or operating system of client computing device **102**, and can include, for example, file system data (e.g., regular files, file tables, mount points, etc.), operating system data (e.g., registries, event logs, etc.), and the like. The e-mail subclient **112B** can include data generated by an e-mail application operating on client computing device **102**, e.g., mailbox information, folder information, emails, attachments, associated database information, and the like. As described above, the subclients can be logical containers, and the data included in the corresponding primary data **112A** and **112B** may or may not be stored contiguously.

(173) The exemplary storage policy **148A** includes backup copy preferences or rule set **160**, disaster recovery copy preferences or rule set **162**, and compliance copy preferences or rule set **164**. Backup copy rule set **160** specifies that it is associated with file system subclient **166** and email subclient **168**. Each of subclients **166** and **168** are associated with the particular client computing device **102**. Backup copy rule set **160** further specifies that the backup operation will be written to disk library **108A** and designates a particular media agent **144A** to convey the data to disk library **108A**. Finally, backup copy rule set **160** specifies that backup copies created according to rule set **160** are scheduled to be generated hourly and are to be retained for 30 days. In some other embodiments, scheduling information is not included in storage policy **148A** and is instead specified by a separate scheduling policy.

(174) Disaster recovery copy rule set **162** is associated with the same two subclients **166** and **168**. However, disaster recovery copy rule set **162** is associated with tape library **108B**, unlike backup copy rule set **160**. Moreover, disaster recovery copy rule set **162** specifies that a different media agent, namely **144B**, will convey data to tape library **108B**. Disaster recovery copies created according to rule set **162** will be retained for 60 days and will be generated daily. Disaster recovery copies generated according to disaster recovery copy rule set **162** can provide protection in the event of a disaster or other catastrophic data loss that would affect the backup copy **116A** maintained on disk library **108A**.

(175) Compliance copy rule set **164** is only associated with the email subclient **168**, and not the file system subclient **166**. Compliance copies generated according to compliance copy rule set **164** will

therefore not include primary data **112A** from the file system subclient **166**. For instance, the organization may be under an obligation to store and maintain copies of email data for a particular period of time (e.g., 10 years) to comply with state or federal regulations, while similar regulations do not apply to file system data. Compliance copy rule set **164** is associated with the same tape library **108B** and media agent **144B** as disaster recovery copy rule set **162**, although a different storage device or media agent could be used in other embodiments. Finally, compliance copy rule set **164** specifies that the copies it governs will be generated quarterly and retained for 10 years.

(176) Secondary Copy Jobs

(177) A logical grouping of secondary copy operations governed by a rule set and being initiated at a point in time may be referred to as a “secondary copy job” (and sometimes may be called a “backup job,” even though it is not necessarily limited to creating only backup copies). Secondary copy jobs may be initiated on demand as well. Steps 1-9 below illustrate three secondary copy jobs based on storage policy **148A**.

(178) Referring to FIG. **1E**, at step 1, storage manager **140** initiates a backup job according to the backup copy rule set **160**, which logically comprises all the secondary copy operations necessary to effectuate rules **160** in storage policy **148A** every hour, including steps 1-4 occurring hourly. For instance, a scheduling service running on storage manager **140** accesses backup copy rule set **160** or a separate scheduling policy associated with client computing device **102** and initiates a backup job on an hourly basis. Thus, at the scheduled time, storage manager **140** sends instructions to client computing device **102** (i.e., to both data agent **142A** and data agent **142B**) to begin the backup job.

(179) At step 2, file system data agent **142A** and email data agent **142B** on client computing device **102** respond to instructions from storage manager **140** by accessing and processing the respective subclient primary data **112A** and **112B** involved in the backup copy operation, which can be found in primary storage device **104**. Because the secondary copy operation is a backup copy operation, the data agent(s) **142A**, **142B** may format the data into a backup format or otherwise process the data suitable for a backup copy.

(180) At step 3, client computing device **102** communicates the processed file system data (e.g., using file system data agent **142A**) and the processed email data (e.g., using email data agent **142B**) to the first media agent **144A** according to backup copy rule set **160**, as directed by storage manager **140**. Storage manager **140** may further keep a record in management database **146** of the association between media agent **144A** and one or more of: client computing device **102**, file system subclient **112A**, file system data agent **142A**, email subclient **112B**, email data agent **142B**, and/or backup copy **116A**.

(181) The target media agent **144A** receives the data-agent-processed data from client computing device **102**, and at step 4 generates and conveys backup copy **116A** to disk library **108A** to be stored as backup copy **116A**, again at the direction of storage manager **140** and according to backup copy rule set **160**. Media agent **144A** can also update its index **153** to include data and/or metadata related to backup copy **116A**, such as information indicating where the backup copy **116A** resides on disk library **108A**, where the email copy resides, where the file system copy resides, data and metadata for cache retrieval, etc. Storage manager **140** may similarly update its index **150** to include information relating to the secondary copy operation, such as information relating to the type of operation, a physical location associated with one or more copies created by the operation, the time the operation was performed, status information relating to the operation, the components involved in the operation, and the like. In some cases, storage manager **140** may update its index **150** to include some or all of the information stored in index **153** of media agent **144A**. At this point, the backup job may be considered complete. After the 30-day retention period expires, storage manager **140** instructs media agent **144A** to delete backup copy **116A** from disk library **108A** and indexes **150** and/or **153** are updated accordingly.

(182) At step 5, storage manager **140** initiates another backup job for a disaster recovery copy according to the disaster recovery rule set **162**. Illustratively this includes steps 5-7 occurring daily for creating disaster recovery copy **116B**. Illustratively, and by way of illustrating the scalable aspects and off-loading principles embedded in system **100**, disaster recovery copy **116B** is based on backup copy

116A and not on primary data **112A** and **112B**.

(183) At step 6, illustratively based on instructions received from storage manager **140** at step 5, the specified media agent **144B** retrieves the most recent backup copy **116A** from disk library **108A**.

(184) At step 7, again at the direction of storage manager **140** and as specified in disaster recovery copy rule set **162**, media agent **144B** uses the retrieved data to create a disaster recovery copy **116B** and store it to tape library **108B**. In some cases, disaster recovery copy **116B** is a direct, mirror copy of backup copy **116A**, and remains in the backup format. In other embodiments, disaster recovery copy **116B** may be further compressed or encrypted, or may be generated in some other manner, such as by using primary data **112A** and **112B** from primary storage device **104** as sources. The disaster recovery copy operation is initiated once a day and disaster recovery copies **116B** are deleted after 60 days; indexes **153** and/or **150** are updated accordingly when/after each information management operation is executed and/or completed. The present backup job may be considered completed.

(185) At step 8, storage manager **140** initiates another backup job according to compliance rule set **164**, which performs steps 8-9 quarterly to create compliance copy **116C**. For instance, storage manager **140** instructs media agent **144B** to create compliance copy **116C** on tape library **108B**, as specified in the compliance copy rule set **164**.

(186) At step 9 in the example, compliance copy **116C** is generated using disaster recovery copy **116B** as the source. This is efficient, because disaster recovery copy resides on the same secondary storage device and thus no network resources are required to move the data. In other embodiments, compliance copy **116C** is instead generated using primary data **112B** corresponding to the email subclient or using backup copy **116A** from disk library **108A** as source data. As specified in the illustrated example, compliance copies **116C** are created quarterly, and are deleted after ten years, and indexes **153** and/or **150** are kept up-to-date accordingly.

(187) Exemplary Applications of Storage Policies—Information Governance Policies and Classification

(188) Again referring to FIG. 1E, storage manager **140** may permit a user to specify aspects of storage policy **148A**. For example, the storage policy can be modified to include information governance policies to define how data should be managed in order to comply with a certain regulation or business objective. The various policies may be stored, for example, in management database **146**. An information governance policy may align with one or more compliance tasks that are imposed by regulations or business requirements. Examples of information governance policies might include a Sarbanes-Oxley policy, a HIPAA policy, an electronic discovery (e-discovery) policy, and so on.

(189) Information governance policies allow administrators to obtain different perspectives on an organization's online and offline data, without the need for a dedicated data silo created solely for each different viewpoint. As described previously, the data storage systems herein build an index that reflects the contents of a distributed data set that spans numerous clients and storage devices, including both primary data and secondary copies, and online and offline copies. An organization may apply multiple information governance policies in a top-down manner over that unified data set and indexing schema in order to view and manipulate the data set through different lenses, each of which is adapted to a particular compliance or business goal. Thus, for example, by applying an e-discovery policy and a Sarbanes-Oxley policy, two different groups of users in an organization can conduct two very different analyses of the same underlying physical set of data/copies, which may be distributed throughout the information management system.

(190) An information governance policy may comprise a classification policy, which defines a taxonomy of classification terms or tags relevant to a compliance task and/or business objective. A classification policy may also associate a defined tag with a classification rule. A classification rule defines a particular combination of criteria, such as users who have created, accessed or modified a document or data object; file or application types; content or metadata keywords; clients or storage locations; dates of data creation and/or access; review status or other status within a workflow (e.g., reviewed or un-reviewed); modification times or types of modifications; and/or any other data attributes in any combination, without limitation. A classification rule may also be defined using other classification tags in the taxonomy. The various criteria used to define a classification rule may be

combined in any suitable fashion, for example, via Boolean operators, to define a complex classification rule. As an example, an e-discovery classification policy might define a classification tag “privileged” that is associated with documents or data objects that (1) were created or modified by legal department staff, or (2) were sent to or received from outside counsel via email, or (3) contain one of the following keywords: “privileged” or “attorney” or “counsel,” or other like terms. Accordingly, all these documents or data objects will be classified as “privileged.”

(191) One specific type of classification tag, which may be added to an index at the time of indexing, is an “entity tag.” An entity tag may be, for example, any content that matches a defined data mask format. Examples of entity tags might include, e.g., social security numbers (e.g., any numerical content matching the formatting mask XXX-XX-XXXX), credit card numbers (e.g., content having a 13-16 digit string of numbers), SKU numbers, product numbers, etc. A user may define a classification policy by indicating criteria, parameters or descriptors of the policy via a graphical user interface, such as a form or page with fields to be filled in, pull-down menus or entries allowing one or more of several options to be selected, buttons, sliders, hypertext links or other known user interface tools for receiving user input, etc. For example, a user may define certain entity tags, such as a particular product number or project ID. In some implementations, the classification policy can be implemented using cloud-based techniques. For example, the storage devices may be cloud storage devices, and the storage manager **140** may execute cloud service provider API over a network to classify data stored on cloud storage devices.

(192) Restore Operations from Secondary Copies

(193) While not shown in FIG. **1E**, at some later point in time, a restore operation can be initiated involving one or more of secondary copies **116A**, **116B**, and **116C**. A restore operation logically takes a selected secondary copy **116**, reverses the effects of the secondary copy operation that created it, and stores the restored data to primary storage where a client computing device **102** may properly access it as primary data. A media agent **144** and an appropriate data agent **142** (e.g., executing on the client computing device **102**) perform the tasks needed to complete a restore operation. For example, data that was encrypted, compressed, and/or deduplicated in the creation of secondary copy **116** will be correspondingly rehydrated (reversing deduplication), uncompressed, and unencrypted into a format appropriate to primary data. Metadata stored within or associated with the secondary copy **116** may be used during the restore operation. In general, restored data should be indistinguishable from other primary data **112**. Preferably, the restored data has fully regained the native format that may make it immediately usable by application **110**.

(194) As one example, a user may manually initiate a restore of backup copy **116A**, e.g., by interacting with user interface **158** of storage manager **140** or with a web-based console with access to system **100**. Storage manager **140** may access data in its index **150** and/or management database **146** (and/or the respective storage policy **148A**) associated with the selected backup copy **116A** to identify the appropriate media agent **144A** and/or secondary storage device **108A** where the secondary copy resides. The user may be presented with a representation (e.g., stub, thumbnail, listing, etc.) and metadata about the selected secondary copy, in order to determine whether this is the appropriate copy to be restored, e.g., date that the original primary data was created. Storage manager **140** will then instruct media agent **144A** and an appropriate data agent **142** on the target client computing device **102** to restore secondary copy **116A** to primary storage device **104**. A media agent may be selected for use in the restore operation based on a load balancing algorithm, an availability based algorithm, or other criteria. The selected media agent, e.g., **144A**, retrieves secondary copy **116A** from disk library **108A**. For instance, media agent **144A** may access its index **153** to identify a location of backup copy **116A** on disk library **108A**, or may access location information residing on disk library **108A** itself.

(195) In some cases a backup copy **116A** that was recently created or accessed, may be cached to speed up the restore operation. In such a case, media agent **144A** accesses a cached version of backup copy **116A** residing in index **153**, without having to access disk library **108A** for some or all of the data. Once it has retrieved backup copy **116A**, the media agent **144A** communicates the data to the requesting client computing device **102**. Upon receipt, file system data agent **142A** and email data agent **142B** may unpack (e.g., restore from a backup format to the native application format) the data

in backup copy **116A** and restore the unpackaged data to primary storage device **104**. In general, secondary copies **116** may be restored to the same volume or folder in primary storage device **104** from which the secondary copy was derived; to another storage location or client computing device **102**; to shared storage, etc. In some cases, the data may be restored so that it may be used by an application **110** of a different version/vintage from the application that created the original primary data **112**.

(196) Exemplary Secondary Copy Formatting

(197) The formatting and structure of secondary copies **116** can vary depending on the embodiment. In some cases, secondary copies **116** are formatted as a series of logical data units or “chunks” (e.g., 512 MB, 1 GB, 2 GB, 4 GB, or 8 GB chunks). This can facilitate efficient communication and writing to secondary storage devices **108**, e.g., according to resource availability. For example, a single secondary copy **116** may be written on a chunk-by-chunk basis to one or more secondary storage devices **108**. In some cases, users can select different chunk sizes, e.g., to improve throughput to tape storage devices. Generally, each chunk can include a header and a payload. The payload can include files (or other data units) or subsets thereof included in the chunk, whereas the chunk header generally includes metadata relating to the chunk, some or all of which may be derived from the payload. For example, during a secondary copy operation, media agent **144**, storage manager **140**, or other component may divide files into chunks and generate headers for each chunk by processing the files. Headers can include a variety of information such as file and/or volume identifier(s), offset(s), and/or other information associated with the payload data items, a chunk sequence number, etc. Importantly, in addition to being stored with secondary copy **116** on secondary storage device **108**, chunk headers can also be stored to index **153** of the associated media agent(s) **144** and/or to index **150** associated with storage manager **140**. This can be useful for providing faster processing of secondary copies **116** during browsing, restores, or other operations. In some cases, once a chunk is successfully transferred to a secondary storage device **108**, the secondary storage device **108** returns an indication of receipt, e.g., to media agent **144** and/or storage manager **140**, which may update their respective indexes **153**, **150** accordingly. During restore, chunks may be processed (e.g., by media agent **144**) according to the information in the chunk header to reassemble the files.

(198) Data can also be communicated within system **100** in data channels that connect client computing devices **102** to secondary storage devices **108**. These data channels can be referred to as “data streams,” and multiple data streams can be employed to parallelize an information management operation, improving data transfer rate, among other advantages. Example data formatting techniques including techniques involving data streaming, chunking, and the use of other data structures in creating secondary copies are described in U.S. Pat. Nos. 7,315,923, 8,156,086, and 8,578,120.

(199) FIGS. 1F and 1G are diagrams of example data streams **170** and **171**, respectively, which may be employed for performing information management operations. Referring to FIG. 1F, data agent **142** forms data stream **170** from source data associated with a client computing device **102** (e.g., primary data **112**). Data stream **170** is composed of multiple pairs of stream header **172** and stream data (or stream payload) **174**. Data streams **170** and **171** shown in the illustrated example are for a single-instanced storage operation, and a stream payload **174** therefore may include both single-instance (SI) data and/or non-SI data. A stream header **172** includes metadata about the stream payload **174**. This metadata may include, for example, a length of the stream payload **174**, an indication of whether the stream payload **174** is encrypted, an indication of whether the stream payload **174** is compressed, an archive file identifier (ID), an indication of whether the stream payload **174** is single instanceable, and an indication of whether the stream payload **174** is a start of a block of data.

(200) Referring to FIG. 1G, data stream **171** has the stream header **172** and stream payload **174** aligned into multiple data blocks. In this example, the data blocks are of size 64 KB. The first two stream header **172** and stream payload **174** pairs comprise a first data block of size 64 KB. The first stream header **172** indicates that the length of the succeeding stream payload **174** is 63 KB and that it is the start of a data block. The next stream header **172** indicates that the succeeding stream payload **174** has a length of 1 KB and that it is not the start of a new data block. Immediately following stream payload **174** is a pair comprising an identifier header **176** and identifier data **178**. The identifier header **176**

includes an indication that the succeeding identifier data **178** includes the identifier for the immediately previous data block. The identifier data **178** includes the identifier that the data agent **142** generated for the data block. The data stream **171** also includes other stream header **172** and stream payload **174** pairs, which may be for SI data and/or non-SI data.

(201) FIG. 1H is a diagram illustrating data structures **180** that may be used to store blocks of SI data and non-SI data on a storage device (e.g., secondary storage device **108**). According to certain embodiments, data structures **180** do not form part of a native file system of the storage device. Data structures **180** include one or more volume folders **182**, one or more chunk folders **184/185** within the volume folder **182**, and multiple files within chunk folder **184**. Each chunk folder **184/185** includes a metadata file **186/187**, a metadata index file **188/189**, one or more container files **190/191/193**, and a container index file **192/194**. Metadata file **186/187** stores non-SI data blocks as well as links to SI data blocks stored in container files. Metadata index file **188/189** stores an index to the data in the metadata file **186/187**. Container files **190/191/193** store SI data blocks. Container index file **192/194** stores an index to container files **190/191/193**. Among other things, container index file **192/194** stores an indication of whether a corresponding block in a container file **190/191/193** is referred to by a link in a metadata file **186/187**. For example, data block B2 in the container file **190** is referred to by a link in metadata file **187** in chunk folder **185**. Accordingly, the corresponding index entry in container index file **192** indicates that data block B2 in container file **190** is referred to. As another example, data block B1 in container file **191** is referred to by a link in metadata file **187**, and so the corresponding index entry in container index file **192** indicates that this data block is referred to.

(202) As an example, data structures **180** illustrated in FIG. 1H may have been created as a result of separate secondary copy operations involving two client computing devices **102**. For example, a first secondary copy operation on a first client computing device **102** could result in the creation of the first chunk folder **184**, and a second secondary copy operation on a second client computing device **102** could result in the creation of the second chunk folder **185**. Container files **190/191** in the first chunk folder **184** would contain the blocks of SI data of the first client computing device **102**. If the two client computing devices **102** have substantially similar data, the second secondary copy operation on the data of the second client computing device **102** would result in media agent **144** storing primarily links to the data blocks of the first client computing device **102** that are already stored in the container files **190/191**. Accordingly, while a first secondary copy operation may result in storing nearly all of the data subject to the operation, subsequent secondary storage operations involving similar data may result in substantial data storage space savings, because links to already stored data blocks can be stored instead of additional instances of data blocks.

(203) If the operating system of the secondary storage computing device **106** on which media agent **144** operates supports sparse files, then when media agent **144** creates container files **190/191/193**, it can create them as sparse files. A sparse file is a type of file that may include empty space (e.g., a sparse file may have real data within it, such as at the beginning of the file and/or at the end of the file, but may also have empty space in it that is not storing actual data, such as a contiguous range of bytes all having a value of zero). Having container files **190/191/193** be sparse files allows media agent **144** to free up space in container files **190/191/193** when blocks of data in container files **190/191/193** no longer need to be stored on the storage devices. In some examples, media agent **144** creates a new container file **190/191/193** when a container file **190/191/193** either includes 100 blocks of data or when the size of the container file **190** exceeds 50 MB. In other examples, media agent **144** creates a new container file **190/191/193** when a container file **190/191/193** satisfies other criteria (e.g., it contains from approx. 100 to approx. 1000 blocks or when its size exceeds approximately 50 MB to 1 GB). In some cases, a file on which a secondary copy operation is performed may comprise a large number of data blocks. For example, a 100 MB file may comprise 400 data blocks of size 256 KB. If such a file is to be stored, its data blocks may span more than one container file, or even more than one chunk folder. As another example, a database file of 20 GB may comprise over 40,000 data blocks of size 512 KB. If such a database file is to be stored, its data blocks will likely span multiple container files, multiple chunk folders, and potentially multiple volume folders. Restoring such files may require accessing multiple container files, chunk folders, and/or volume folders to obtain the requisite data

blocks.

(204) Using Backup Data for Replication and Disaster Recovery (“Live Synchronization”)

(205) There is an increased demand to off-load resource intensive information management tasks (e.g., data replication tasks) away from production devices (e.g., physical or virtual client computing devices) in order to maximize production efficiency. At the same time, enterprises expect access to readily-available up-to-date recovery copies in the event of failure, with little or no production downtime.

(206) FIG. 2A illustrates a system **200** configured to address these and other issues by using backup or other secondary copy data to synchronize a source subsystem **201** (e.g., a production site) with a destination subsystem **203** (e.g., a failover site). Such a technique can be referred to as “live synchronization” and/or “live synchronization replication.” In the illustrated embodiment, the source client computing devices **202a** include one or more virtual machines (or “VMs”) executing on one or more corresponding VM host computers **205a**, though the source need not be virtualized. The destination site **203** may be at a location that is remote from the production site **201**, or may be located in the same data center, without limitation. One or more of the production site **201** and destination site **203** may reside at data centers at known geographic locations, or alternatively may operate “in the cloud.”

(207) The synchronization can be achieved by generally applying an ongoing stream of incremental backups from the source subsystem **201** to the destination subsystem **203**, such as according to what can be referred to as an “incremental forever” approach. FIG. 2A illustrates an embodiment of a data flow which may be orchestrated at the direction of one or more storage managers (not shown). At step 1, the source data agent(s) **242a** and source media agent(s) **244a** work together to write backup or other secondary copies of the primary data generated by the source client computing devices **202a** into the source secondary storage device(s) **208a**. At step 2, the backup/secondary copies are retrieved by the source media agent(s) **244a** from secondary storage. At step 3, source media agent(s) **244a** communicate the backup/secondary copies across a network to the destination media agent(s) **244b** in destination subsystem **203**.

(208) As shown, the data can be copied from source to destination in an incremental fashion, such that only changed blocks are transmitted, and in some cases multiple incremental backups are consolidated at the source so that only the most current changed blocks are transmitted to and applied at the destination. An example of live synchronization of virtual machines using the “incremental forever” approach is found in U.S. Patent Application No. 62/265,339 entitled “Live Synchronization and Management of Virtual Machines across Computing and Virtualization Platforms and Using Live Synchronization to Support Disaster Recovery.” Moreover, a deduplicated copy can be employed to further reduce network traffic from source to destination. For instance, the system can utilize the deduplicated copy techniques described in U.S. Pat. No. 9,239,687, entitled “Systems and Methods for Retaining and Using Data Block Signatures in Data Protection Operations.”

(209) At step 4, destination media agent(s) **244b** write the received backup/secondary copy data to the destination secondary storage device(s) **208b**. At step 5, the synchronization is completed when the destination media agent(s) and destination data agent(s) **242b** restore the backup/secondary copy data to the destination client computing device(s) **202b**. The destination client computing device(s) **202b** may be kept “warm” awaiting activation in case failure is detected at the source. This synchronization/replication process can incorporate the techniques described in U.S. patent application Ser. No. 14/721,971, entitled “Replication Using Deduplicated Secondary Copy Data.”

(210) Where the incremental backups are applied on a frequent, on-going basis, the synchronized copies can be viewed as mirror or replication copies. Moreover, by applying the incremental backups to the destination site **203** using backup or other secondary copy data, the production site **201** is not burdened with the synchronization operations. Because the destination site **203** can be maintained in a synchronized “warm” state, the downtime for switching over from the production site **201** to the destination site **203** is substantially less than with a typical restore from secondary storage. Thus, the production site **201** may flexibly and efficiently fail over, with minimal downtime and with relatively up-to-date data, to a destination site **203**, such as a cloud-based failover site. The destination site **203**

can later be reverse synchronized back to the production site **201**, such as after repairs have been implemented or after the failure has passed.

(211) Integrating With the Cloud Using File System Protocols

(212) Given the ubiquity of cloud computing, it can be increasingly useful to provide data protection and other information management services in a scalable, transparent, and highly plug-able fashion. FIG. 2B illustrates an information management system **200** having an architecture that provides such advantages, and incorporates use of a standard file system protocol between primary and secondary storage subsystems **217**, **218**. As shown, the use of the network file system (NFS) protocol (or any another appropriate file system protocol such as that of the Common Internet File System (CIFS)) allows data agent **242** to be moved from the primary storage subsystem **217** to the secondary storage subsystem **218**. For instance, as indicated by the dashed box **206** around data agent **242** and media agent **244**, data agent **242** can co-reside with media agent **244** on the same server (e.g., a secondary storage computing device such as component **106**), or in some other location in secondary storage subsystem **218**.

(213) Where NFS is used, for example, secondary storage subsystem **218** allocates an NFS network path to the client computing device **202** or to one or more target applications **210** running on client computing device **202**. During a backup or other secondary copy operation, the client computing device **202** mounts the designated NFS path and writes data to that NFS path. The NFS path may be obtained from NFS path data **215** stored locally at the client computing device **202**, and which may be a copy of or otherwise derived from NFS path data **219** stored in the secondary storage subsystem **218**.

(214) Write requests issued by client computing device(s) **202** are received by data agent **242** in secondary storage subsystem **218**, which translates the requests and works in conjunction with media agent **244** to process and write data to a secondary storage device(s) **208**, thereby creating a backup or other secondary copy. Storage manager **240** can include a pseudo-client manager **217**, which coordinates the process by, among other things, communicating information relating to client computing device **202** and application **210** (e.g., application type, client computing device identifier, etc.) to data agent **242**, obtaining appropriate NFS path data from the data agent **242** (e.g., NFS path information), and delivering such data to client computing device **202**.

(215) Conversely, during a restore or recovery operation client computing device **202** reads from the designated NFS network path, and the read request is translated by data agent **242**. The data agent **242** then works with media agent **244** to retrieve, re-process (e.g., re-hydrate, decompress, decrypt), and forward the requested data to client computing device **202** using NFS.

(216) By moving specialized software associated with system **200** such as data agent **242** off the client computing devices **202**, the illustrative architecture effectively decouples the client computing devices **202** from the installed components of system **200**, improving both scalability and plug-ability of system **200**. Indeed, the secondary storage subsystem **218** in such environments can be treated simply as a read/write NFS target for primary storage subsystem **217**, without the need for information management software to be installed on client computing devices **202**. As one example, an enterprise implementing a cloud production computing environment can add VM client computing devices **202** without installing and configuring specialized information management software on these VMs.

Rather, backups and restores are achieved transparently, where the new VMs simply write to and read from the designated NFS path. An example of integrating with the cloud using file system protocols or so-called “infinite backup” using NFS share is found in U.S. Patent Application No. 62/294,920, entitled “Data Protection Operations Based on Network Path Information.” Examples of improved data restoration scenarios based on network-path information, including using stored backups effectively as primary data sources, may be found in U.S. Patent Application No. 62/297,057, entitled “Data Restoration Operations Based on Network Path Information.”

(217) Highly Scalable Managed Data Pool Architecture

(218) Enterprises are seeing explosive data growth in recent years, often from various applications running in geographically distributed locations. FIG. 2C shows a block diagram of an example of a highly scalable, managed data pool architecture useful in accommodating such data growth. The illustrated system **200**, which may be referred to as a “web-scale” architecture according to certain

embodiments, can be readily incorporated into both open compute/storage and common-cloud architectures.

(219) The illustrated system **200** includes a grid **245** of media agents **244** logically organized into a control tier **231** and a secondary or storage tier **233**. Media agents assigned to the storage tier **233** can be configured to manage a secondary storage pool **208** as a deduplication store, and be configured to receive client write and read requests from the primary storage subsystem **217**, and direct those requests to the secondary tier **233** for servicing. For instance, media agents CMA1-CMA3 in the control tier **231** maintain and consult one or more deduplication databases **247**, which can include deduplication information (e.g., data block hashes, data block links, file containers for deduplicated files, etc.) sufficient to read deduplicated files from secondary storage pool **208** and write deduplicated files to secondary storage pool **208**. For instance, system **200** can incorporate any of the deduplication systems and methods shown and described in U.S. Pat. No. 9,020,900, entitled “Distributed Deduplicated Storage System,” and U.S. Pat. Pub. No. 2014/0201170, entitled “High Availability Distributed Deduplicated Storage System.”

(220) Media agents SMA1-SMA6 assigned to the secondary tier **233** receive write and read requests from media agents CMA1-CMA3 in control tier **231**, and access secondary storage pool **208** to service those requests. Media agents CMA1-CMA3 in control tier **231** can also communicate with secondary storage pool **208**, and may execute read and write requests themselves (e.g., in response to requests from other control media agents CMA1-CMA3) in addition to issuing requests to media agents in secondary tier **233**. Moreover, while shown as separate from the secondary storage pool **208**, deduplication database(s) **247** can in some cases reside in storage devices in secondary storage pool **208**.

(221) As shown, each of the media agents **244** (e.g., CMA1-CMA3, SMA1-SMA6, etc.) in grid **245** can be allocated a corresponding dedicated partition **251A-251I**, respectively, in secondary storage pool **208**. Each partition **251** can include a first portion **253** containing data associated with (e.g., stored by) media agent **244** corresponding to the respective partition **251**. System **200** can also implement a desired level of replication, thereby providing redundancy in the event of a failure of a media agent **244** in grid **245**. Along these lines, each partition **251** can further include a second portion **255** storing one or more replication copies of the data associated with one or more other media agents **244** in the grid.

(222) System **200** can also be configured to allow for seamless addition of media agents **244** to grid **245** via automatic configuration. As one illustrative example, a storage manager (not shown) or other appropriate component may determine that it is appropriate to add an additional node to control tier **231**, and perform some or all of the following: (i) assess the capabilities of a newly added or otherwise available computing device as satisfying a minimum criteria to be configured as or hosting a media agent in control tier **231**; (ii) confirm that a sufficient amount of the appropriate type of storage exists to support an additional node in control tier **231** (e.g., enough disk drive capacity exists in storage pool **208** to support an additional deduplication database **247**); (iii) install appropriate media agent software on the computing device and configure the computing device according to a pre-determined template; (iv) establish a partition **251** in the storage pool **208** dedicated to the newly established media agent **244**; and (v) build any appropriate data structures (e.g., an instance of deduplication database **247**). An example of highly scalable managed data pool architecture or so-called web-scale architecture for storage and data management is found in U.S. Patent Application No. 62/273,286 entitled “Redundant and Robust Distributed Deduplication Data Storage System.”

(223) The embodiments and components thereof disclosed in FIGS. 2A, 2B, and 2C, as well as those in FIGS. 1A-1H, may be implemented in any combination and permutation to satisfy data storage management and information management needs at one or more locations and/or data centers.

(224) Preparing Containerized Applications for Backup Using a Backup Services Container

(225) The illustrative backup services container (e.g., **301**) interoperates with a proprietary data storage management system (e.g., **302**) and/or with third-party backup systems (e.g., **303**). The former configuration provides numerous features and advantages over the latter configuration. Furthermore, some embodiments include one or more components of the proprietary data storage management

with the illustrative backup services container. Some embodiments include one or more components of the proprietary data storage management system in a backup services pod (e.g., **486**) configured in a Kubernetes node. All configurations and embodiments are suitable for cloud and non-cloud computing environments, without limitation.

(226) Cloud Computing. The National Institute of Standards and Technology (NIST) provides the following definition of Cloud Computing characteristics, service models, and deployment models:

(227) Cloud Computing

(228) Cloud computing is a model for enabling ubiquitous, convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction. This cloud model is composed of five essential characteristics, three service models, and four deployment models.

Essential Characteristics: On-demand self-service. A consumer can unilaterally provision computing capabilities, such as server time and network storage, as needed automatically without requiring human interaction with each service provider. Broad network access. Capabilities are available over the network and accessed through standard mechanisms that promote use by heterogeneous thin or thick client platforms (e.g., mobile phones, tablets, laptops, and workstations). Resource pooling. The provider's computing resources are pooled to serve multiple consumers using a multi-tenant model, with different physical and virtual resources dynamically assigned and reassigned according to consumer demand. There is a sense of location independence in that the customer generally has no control or knowledge over the exact location of the provided resources but may be able to specify location at a higher level of abstraction (e.g., country, state, or datacenter). Examples of resources include storage, processing, memory, and network bandwidth. Rapid elasticity. Capabilities can be elastically provisioned and released, in some cases automatically, to scale rapidly outward and inward commensurate with demand. To the consumer, the capabilities available for provisioning often appear to be unlimited and can be appropriated in any quantity at any time. Measured service. Cloud systems automatically control and optimize resource use by leveraging a metering capability.^{sup.1} at some level of abstraction appropriate to the type of service (e.g., storage, processing, bandwidth, and active user accounts). Resource usage can be monitored, controlled, and reported, providing transparency for both the provider and consumer of the utilized service. ^{sup.1} Typically this is done on a pay-per-use or charge-per-use basis.

Service Models: Software as a Service (SaaS). The capability provided to the consumer is to use the provider's applications running on a cloud infrastructure.^{sup.2} The applications are accessible from various client devices through either a thin client interface, such as a web browser (e.g., web-based email), or a program interface. The consumer does not manage or control the underlying cloud infrastructure including network, servers, operating systems, storage, or even individual application capabilities, with the possible exception of limited user-specific application configuration settings. ^{sup.2} A cloud infrastructure is the collection of hardware and software that enables the five essential characteristics of cloud computing. The cloud infrastructure can be viewed as containing both a physical layer and an abstraction layer. The physical layer consists of the hardware resources that are necessary to support the cloud services being provided, and typically includes server, storage and network components. The abstraction layer consists of the software deployed across the physical layer, which manifests the essential cloud characteristics. Conceptually the abstraction layer sits above the physical layer. Platform as a Service (PaaS). The capability provided to the consumer is to deploy onto the cloud infrastructure consumer-created or acquired applications created using programming languages, libraries, services, and tools supported by the provider..^{sup.3} The consumer does not manage or control the underlying cloud infrastructure including network, servers, operating systems, or storage, but has control over the deployed applications and possibly configuration settings for the application-hosting environment. ^{sup.3} This capability does not necessarily preclude the use of compatible programming languages, libraries, services, and tools from other sources. Infrastructure as a Service (IaaS). The capability provided to the consumer is to provision processing, storage, networks, and other fundamental computing resources where the consumer is able to deploy and run

arbitrary software, which can include operating systems and applications. The consumer does not manage or control the underlying cloud infrastructure but has control over operating systems, storage, and deployed applications; and possibly limited control of select networking components (e.g., host firewalls).

Deployment Models: Private cloud. The cloud infrastructure is provisioned for exclusive use by a single organization comprising multiple consumers (e.g., business units). It may be owned, managed, and operated by the organization, a third party, or some combination of them, and it may exist on or off premises. Community cloud. The cloud infrastructure is provisioned for exclusive use by a specific community of consumers from organizations that have shared concerns (e.g., mission, security requirements, policy, and compliance considerations). It may be owned, managed, and operated by one or more of the organizations in the community, a third party, or some combination of them, and it may exist on or off premises. Public cloud. The cloud infrastructure is provisioned for open use by the general public. It may be owned, managed, and operated by a business, academic, or government organization, or some combination of them. It exists on the premises of the cloud provider. Hybrid cloud. The cloud infrastructure is a composition of two or more distinct cloud infrastructures (private, community, or public) that remain unique entities, but are bound together by standardized or proprietary technology that enables data and application portability (e.g., cloud bursting for load balancing between clouds).

(229) Source: Peter Mell, Timothy Grance (September 2011). *The NIST Definition of Cloud Computing*, National Institute of Standards and Technology: U.S. Department of Commerce. Special publication 800-145. nvlpubs.nist.gov/nistpubs/Legacy/SP/nistspecialpublication800-145.pdf (accessed 26 Apr. 2019). Cloud computing aims to allow those who consume the services (whether individuals or organizations) to benefit from the available technologies without the need for deep knowledge about or expertise with each of them. Wikipedia, Cloud Computing, en.wikipedia.org/wiki/Cloud_computing (accessed 26 Apr. 2019). “Cloud computing metaphor: the group of networked elements providing services need not be individually addressed or managed by users; instead, the entire provider-managed suite of hardware and software can be thought of as an amorphous cloud.” Id.

(230) Cloud Service Accounts and Variability in Cloud Services. Cloud service providers such as Amazon, Microsoft, Alibaba, Google, Salesforce, Cisco, etc. provide access to their particular cloud services via cloud service accounts, such as corporate accounts, departmental accounts, individual user accounts, etc. Each cloud service account typically has authentication features, e.g., passwords, certificates, etc., to restrict and control access to the cloud service. Each account also might have service level guarantees and/or other terms and conditions between the cloud service provider and the service subscriber, e.g., a company, a government agency, an individual user. A subscribing entity might have multiple accounts with a cloud service provider, such as an account for the Engineering department, an account for the Finance department, an account for the Human Resources department, other accounts for individual company users, etc., without limitation. Each cloud service account carries different authentication, even though the services subscriber is the same entity.

(231) Different cloud service accounts might differ not just in service level guarantees, but might include different services. For example, one account might include long-term storage resources, whereas another account might be limited to ordinary data storage. For example, some accounts might have access to data processing functions supplied by the cloud service provider, such as machine learning algorithms, statistical analysis packages, etc., whereas other accounts might lack such features. Accordingly, the resources available to the user(s) of cloud service accounts can vary as between accounts, even if the accounts have the same subscriber and the same cloud service provider. Thus, the user experience and the technologies available as between cloud service accounts can vary significantly. Thus, when considering cloud computing, the specifics of cloud service accounts can play a role in the availability and/or portability of resources. Crossing account boundaries can pose technological barriers when considering migration of applications and their cloud services assets.

(232) Cloud Availability Zones. “Availability zones (AZs) are isolated locations within . . . regions from which public cloud services originate and operate. Regions are geographic locations in which

public cloud service providers' data centers reside. Businesses choose one or multiple worldwide availability zones for their services depending on business needs. Businesses select availability zones for a variety of reasons, including compliance and proximity to end customers. Cloud administrators can also choose to replicate services across multiple availability zones to decrease latency or protect resources. Admins can move resources to another availability zone in the event of an outage. Certain cloud services may also be limited to particular regions or AZs.” Source: Margaret Rouse, Definition of Availability Zones, TechTarget, searchaws.techtarget.com/definition/availability-zones (accessed 26 Apr. 2019).

(233) Here is a vendor-specific example of how cloud service availability zones are organized in the Google Cloud: “Certain [Google] Compute Engine resources live in regions or zones. A region is a specific geographical location where you can run your resources. Each region has one or more zones; most regions have three or more zones. For example, the us-central1 region denotes a region in the Central United States that has zones us-central1-a, us-central1-b, us-central1-c, and us-central1-f. Resources that live in a zone, such as instances or persistent disks, are referred to as zonal resources. Other resources, like static external IP addresses, are regional. Regional resources can be used by any resources in that region, regardless of zone, while zonal resources can only be used by other resources in the same zone. For example, disks and instances are both zonal resources. To attach a disk to an instance, both resources must be in the same zone. Similarly, if you want to assign a static IP address to an instance, the instance must be in the same region as the static IP. Only certain resources are region- or zone-specific. Other resources, such as images, are global resources that can be used by any other resources across any location. For information on global, regional, and zonal Compute Engine resources, see Global, Regional, and Zonal Resources.” Source: Google Cloud Regions and Zones, cloud.google.com/compute/docs/regions-zones/ (accessed 26 Apr. 2019) (emphasis added).

(234) Accordingly, when considering cloud computing, availability zones can play a role in the availability and/or portability of resources. Crossing zone boundaries can pose technological barriers when considering migration of applications and their cloud service assets, even when the different availability zones are supplied by the same cloud service provider.

(235) Traditional Non-Cloud (“On-Premises”) Data Centers are Distinguishable from Cloud Computing. Traditional data centers generally do not have cloud computing characteristics. For example, the user experience is generally different, for example in regard to the name space(s) used for the storage, computing, and network resources. Moreover, substantial increases in resources needed by a user are not provisioned on demand. A traditional data center is physically located within the enterprise/organization that owns it. A traditional non-cloud data center might comprise computing resources such as servers, mainframes, virtual servers/clusters, etc.; and/or data storage resources, such as network-attached storage, storage area networks, tape libraries, etc. The owner of the traditional data center procures hardware, software, and network infrastructure (including making the associated capital investments); and manages going-forward planning for the data center. A traditional data center is staffed by professional Information Technology (IT) personnel, who are responsible for the data center's configuration, operation, upgrades, and maintenance. Thus, a traditional non-cloud data center can be thought of as self-managed by its owner/operator for the benefit of in-house users, as compared to cloud computing, which is managed by the cloud service provider and supplied as a service to outside subscribers. Clearly, a cloud computing service also has hardware, software, and networking infrastructure and professionals staffing it, as well as having an owner responsible for housing and paying for the infrastructure. However, the cloud computing service is consumed differently, served differently, and deployed differently compared to non-cloud data centers. Traditional non-cloud data centers are sometimes referred to as “on-premises” data centers, because their facilities are literally within the bounds of the organization that owns the data center. Cloud service providers' data centers generally are not within the bounds of the subscriber organization and are consumed “at a distance” “in the cloud.”

(236) Accordingly, when considering cloud computing versus non-cloud data center deployment, the choice can play a role in the availability and/or portability of resources. Crossing boundaries between non-cloud data centers and cloud computing can pose technological barriers. For example, storing a

database at a non-cloud data center might require different resources and/or access features/controls than storing the database at a cloud computing service. Thus, moving the database from the non-cloud data center to a cloud service account may require data conversion, re-configuration, and/or adaptation that go above and beyond merely copying the database. Likewise for virtual machines (VMs). Conversely, moving data, applications, VMs, and/or web services from cloud computing to a non-cloud data center also can involve data conversion, re-configuration, and/or adaptation to ensure success.

(237) Service Models. Differences in service models, comparing non-cloud “on-premises” data centers versus IaaS versus PaaS versus SaaS, can yield different performance and cost profiles. Different service models can affect resource availability and/or portability of distributed/serverless applications, at least because the management of different resources rests with different providers and governed by different terms and conditions. See, e.g., Stephen Watts, SaaS vs PaaS vs IaaS: What's The Difference and How To Choose, BMC Blogs, BMC Software, Inc., www.bmc.com/blogs/saas-vs-paas-vs-iaas-whats-the-difference-and-how-to-choose/ (accessed 26 Apr. 2019).

(238) In regard to the figures described herein, other embodiments are possible within the scope of the present invention, such that the above-recited components, steps, blocks, operations, messages, requests, queries, and/or instructions are differently arranged, sequenced, sub-divided, organized, and/or combined. In some embodiments, a different component may initiate or execute a given operation.

(239) FIG. 3A is a block diagram illustrating some salient portions of a system **300** comprising backup services container **301** and data storage management system **302**, according to an illustrative embodiment. FIG. 3A depicts: data storage management system **302**; and container-orchestration (Kubernetes) pod **310**, comprising backup services container **301**, container **319** comprising containerized applications **320**, and data storage volumes **330**. In some embodiments, pod **310** executes in a cloud computing system that provides a cloud computing environment. In some embodiments, pod **310** executes in a non-cloud data center. In some embodiments, some or all components of data storage management system **302** execute in a cloud computing environment provided by a cloud computing system, e.g., Microsoft Azure, Amazon Web Services, Google Cloud, etc., without limitation. In other embodiments, some or all components of data storage management system **302** execute in a non-cloud computing environment.

(240) System **300** is defined as a combination of backup services container **301** and data storage management system **302**. The other contents pod **310** are not necessarily part of system **300**.

(241) Backup services container **301** is based on an operating system-virtualization (OS-virtualization) service and is illustratively embodied as a Docker container. See, e.g., [http://en.wikipedia.org/wiki/Docker_\(software\)](http://en.wikipedia.org/wiki/Docker_(software)) (accessed Jul. 5, 2019). OS-virtualization enables containers to be isolated from each other, each container being run by a single operating system kernel. Container contents include software, libraries, and configuration files that communicate with each other. In alternative embodiments, backup services container **301** is based on an OS-virtualization service, but is not a Docker container, without limitation. Containers are generally known in the art, but the novel contents and functionality configured in container **301** will be appreciated by those having ordinary skill in the art after reading the present disclosure. More details are given in other figures.

(242) Data storage management system **302** is analogous to system **100** described in more detail elsewhere herein, and additionally comprises functionality for interoperating with backup services container **301** and for offering enhanced features resulting from this interoperability. Data storage management system **302** is a proprietary system provided by the present applicant, Commvault Systems, Inc., of Tinton Falls, New Jersey, USA. More details are given in other figures.

(243) Container-orchestration pod **310** is illustratively embodied as a Kubernetes pod that operates within a Kubernetes container-orchestration system. See, e.g., <http://kubernetes.io/docs/tutorials/kubernetes-basics/explore/explore-intro/> (accessed Jul. 3, 2019); see also <http://en.wikipedia.org/wiki/Kubernetes>. “Kubernetes (K8s) is an open-source system for automating deployment, scaling, and management of containerized applications. It groups containers

that make up an application into logical units for easy management and discovery.” Kubernetes home page <http://kubernetes.io/> (accessed Jul. 10, 2019). See also <http://kubernetes.io/docs/concepts/overview/what-is-kubernetes/> (accessed Jul. 10, 2019).

(244) Pod **310** groups containers whose components are co-located on the host machine and share resources. Within the pod, containers reference each other locally. A pod defines one or more data storage volumes, such as a network disk or local disk directory, which is exposed to the containers in the pod. In alternative embodiments, pod **310** is based on a container-orchestration system other than Kubernetes, without limitation. Illustratively, pod **310** comprises container **301** and container **319** and storage volumes **330-1** and **330-2**, though the invention is not limited to any particular number of containers and/or storage volumes. More details are given in other figures.

(245) Container **319** is based on an operating system-virtualization (OS-virtualization) service and is illustratively embodied by a Docker container, like backup services container **301**. In alternative embodiments, container **319** is based on an OS-virtualization service that is not a Docker container, without limitation. Application containerization is an OS-level virtualization method used to deploy and run distributed applications without launching an entire virtual machine (VM) for each app [application]. Multiple isolated applications or services run on a single host and access the same OS kernel. Containers work on bare-metal systems, cloud instances and virtual machines, across Linux and select Windows and Mac OSes . . . Application containers include the runtime components—such as files, environment variables and libraries—necessary to run the desired software. Application containers consume fewer resources than a comparable deployment on virtual machines because containers share resources without a full operating system to underpin each app. The complete set of information to execute in a container is the image. The container engine deploys these images on hosts. The most common app [application] containerization technology is Docker, specifically the open source Docker Engine and containers based on universal runtime runC.”

<http://searchitoperations.techtarget.com/definition/application-containerization-app-containerization> (accessed Jul. 5, 2019).

(246) Container **319** illustratively comprises three applications **320**, which are referred to as “containerized applications,” because they operate in a container. There is no limit on how many and what kind of containerized applications **320** are configured in a container **319**. Likewise, there is no limit on how many containers **319** are configured in a pod **310**.

(247) Containerized applications **320** (e.g., MySQL DBMS **320-1**, PostgreSQL DBMS **320-2**, Microsoft SQL DBMS **320-3**, etc., without limitation) are applications that are configured to execute within container **319** within pod **310**. Each application comprises executable software that executes within a running container, such as container **319**. Although the depicted containerized applications **320** are all DBMSs, containerization in general and the present invention in particular are not limited to DBMS applications. Other exemplary containerized applications **320** include file managers for accessing file systems, web servers, etc., without limitation. Any application can be containerized.

(248) Data storage volumes **330** are embodied as Kubernetes volumes when implemented in a Kubernetes pod. “A Kubernetes Volume provides persistent storage that exists for the lifetime of the pod itself. This storage can also be used as shared disk space for containers within the pod. Volumes are mounted at specific mount points within the container, which are defined by the pod configuration.” <http://en.wikipedia.org/wiki/Kubernetes#Volumes> (accessed Jul. 5, 2019). Although the present figure depicts two volumes **330** configured in pod **310**, the invention is not so limited, and there is no limit on how many data storage volumes **330** are configured in given pod **310**.

(249) Although the present figure depicts backup services container **301** in communication with each of the depicted containerized applications **320**, the communicative couplings may be indirect, i.e., via container-to-container interfaces, and the depicted arrows are to be interpreted as logical views.

Furthermore, the communicative couplings are generally not ongoing or persistent, but instead they occur on demand. For example, backup services container **301** runs discovery logic as requested by a storage manager in data storage management system **302**; or runs discovery logic and/or selection logic when receiving notice of a backup operation. Accordingly, after executing the discovery logic and/or selection logic, backup services container **301** discontinues its communicative coupling to the

target containerized applications **320**. See also FIG. 5.

(250) FIG. 3B is a block diagram illustrating backup services container **301** interoperating with a third-party backup system **303**, according to another illustrative embodiment. FIG. 3B is similar to FIG. 3A, except that it does not show the outline of container **319** to simplify the rendering of the present figure. FIG. 3B is further similar to FIG. 3A, except that instead of communicating with proprietary data storage management system **302**, backup services container **301** communicates with a third-party backup system **303**. In some embodiments, pod **310** executes in a cloud computing system that provides a cloud computing environment. In some embodiments, pod **310** executes in a non-cloud data center. In some embodiments, some or all components of backup system **303** execute in a cloud computing environment provided by a cloud computing system, e.g., Microsoft Azure, Amazon Web Services, Google Cloud, etc., without limitation. In other embodiments, some or all components of backup system **303** execute in a non-cloud computing environment.

(251) Backup system **303** is a system that can perform backup operations for data generated by containerized applications **320** in pod **310**. An example third-party backup system is Velero (see, e.g., <http://github.com/heptio/velero>) which includes open source utilities for backing up Kubernetes applications. Backup systems **303** are well known in the art. Backup services container **301** is configured to communicate with any third-party backup system **303**, e.g., receiving notice of a backup operation; reporting that it has prepared certain pod contents for backup (e.g., one or more containerized applications **320**); and receiving notice that the backup operation has completed (e.g., snapshot has been taken).

(252) FIG. 3C is a block diagram of an illustrative overview of Kubernetes pods as described in the prior art. An original diagram was obtained from <http://kubernetes.io/docs/tutorials/kubernetes-basics/explore/explore-intro/> (accessed Jul. 3, 2019) and reference numbers were added to enhance the reader's understanding of the present disclosure. "A Pod is a group of one or more application containers (such as Docker or rkt) and includes shared storage (volumes), IP address and information about how to run them . . . such as the container image version or specific ports to use." Id. (accessed May 22, 2020).

(253) All the depicted components in the present figure are well known in the art, e.g., pod IP addresses **312**, Kubernetes pods **311** (**311-1**, **311-2**, **311-3**, **311-4**), containerized applications **320**, and volumes **330**. As evidenced by the diagram, a pod may comprise one or more containerized applications **320** and/or one or more data storage volumes **330**, and is associated with a unique pod IP address **312**.

(254) A notable distinction here is that prior art pods **311** do not comprise the illustrative backup services container **301**, which is configured into container-orchestration (Kubernetes) pods **310**, as shown in other figures herein. For comparison, see, e.g., FIG. 4C.

(255) FIG. 3D is a block diagram of an illustrative overview of a Kubernetes node **313** comprising a plurality of Kubernetes pods **311** as described in the prior art. An original diagram was obtained from <http://kubernetes.io/docs/tutorials/kubernetes-basics/explore/explore-intro/> (accessed Jul. 3, 2019) and reference numbers were added to enhance the reader's understanding of the present disclosure. "A node is a worker machine in Kubernetes and may be a VM [virtual machine] or physical machine, depending on the cluster. Multiple Pods can run on one Node." Id. (accessed May 22, 2020). "The kubelet is the primary 'node agent' that runs on each node. It can register the node with the apiserver using one of: the hostname; a flag to override the hostname; or specific logic for a cloud provider . . . The kubelet takes a set of PodSpecs [object that describes a pod] that are provided through various mechanisms (primarily through the apiserver) and ensures that the containers described in those PodSpecs are running and healthy. The kubelet doesn't manage containers which were not created by Kubernetes." <http://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/> (accessed May 22, 2020).

(256) All the depicted components in the present figure are well known in the art, e.g., pod IP addresses **312**, Kubernetes node **313**, node processes (Kubelet) **314**, Kubernetes pods **311** (**311-1**, **311-2**, **311-3**, **311-4**), containerized applications **320**, and volumes **330**. As in FIG. 3C, a notable distinction here is that pods **311** do not comprise the illustrative backup services container **301**, which is

configured into container-orchestration (Kubernetes) pods **310**. For comparison, see, e.g., FIG. 4C. (257) FIG. 4A is a block diagram depicting some salient details of backup services container **301** and data storage management system **302**. FIG. 4A depicts: backup services container **301** comprising backup preparation toolkits **401**, discovery logic **450**, selection logic **460**, and backup interface logic **470**; data storage management system **302** comprising secondary copies **116**, backup proxy **406**, storage manager **440**, and management database **446**; and containerized applications **320**. Illustratively, backup services container **301** communicates with storage manager **440** using interface logic **470** as depicted by the bi-directional arrow therebetween. Again, one or more containers **319**, which host containerized application(s) **320**, are not shown by outlines in the present figure to simplify the presentation.

(258) Secondary copies **116** are described in more detail elsewhere herein. In the present context, each copy **116** is generated by a secondary copy operation that takes as its source primary data generated by a containerized application **320** (or a snapshot of the primary data). The source primary data and/or the snapshot thereof is located in the container hosting the application **320**, in another data storage resource within the container, in another container, in other cloud-based data storage, or in non-cloud data storage, without limitation. Likewise, there is no limitation on where the secondary copies **116** are stored.

(259) Backup preparation toolkits **401** (e.g., **401-a** . . . **401-f**, etc.) are groupings of scripts and/or commands accompanied by an execution environment therefor, e.g., runtime C, python, etc. Each toolkit is configured for a particular type and/or specific version of an application that requires preparation for backup, e.g., DBMSs, file systems, file managers for accessing a file system, etc. Different versions of the same kind of DBMS (e.g., MySQL 1.0 vs. MySQL 5.6) might use different commands or APIs for quiescing/unquiescing, and therefore multiple backup toolkits might be required to support these differences. Likewise, different file systems (or file managers for accessing the file systems) also have specific commands that are not interchangeable with other file systems and may change from one release to another, thus requiring distinct backup toolkits. Thus, each toolkit **401** comprises resources sufficient to prepare a certain kind and version(s) of an application for backup. For example, if certain scripts are written in C, a runtime C execution environment is also included in the backup toolkit **401** to ensure that the scripts can properly execute. Likewise, if certain commands are written in python, a python execution environment is needed for and is included in the particular toolkit **401** to run the commands against the targeted containerized application **320**.

(260) According to the illustrative embodiments, each backup services container **301** comprises as many backup preparation toolkits **401** as feasible without regard to which applications **320** might be present in containers **319**. In fact, knowledge of those applications will be obtained after backup services container **301** is added to Kubernetes pod **310** (e.g., using discovery logic **450**). Therefore, one of the key advantages of the present embodiments is that each backup services container **301** is “maximally” equipped with a wide range of toolkits **401** that may or may not be tapped within the present Kubernetes pod **310**. Thus, backup services container **301** brings wide applicability to a range of containerized applications **320** that might be configured in the present pod **310** or in other pods **310**. Although FIG. 4A depicts only six backup toolkits **401**, there is no limitation on how many and how many types of toolkits **401** are configured in a backup services container **301**. Likewise, FIG. 4A depicts only three target containerized applications **320**, each one of a different type than the others, but the invention has no limitations on numbers and/or types of containerized applications **320** that can be targeted for backup preparation by a backup services container **301**.

(261) Backup proxy **406** comprises a computing device of any kind (e.g., virtualized, non-virtualized, cloud-based, non-cloud, etc.) that hosts backup components (e.g., data agents, media agents) for generating secondary copies **116**. Illustratively, any combination of data agents **142/442** and media agents **144/444** can be configured in backup proxy **406**. In some embodiments, data agents **142/442** and media agents **144/444** execute on distinct computing devices, but they are shown here on the same computing device as a logical collection of backup resources collectively referred to as backup proxy **406**.

(262) Storage manager **440** is analogous to storage manager **140** and additionally comprises features

for interoperating with backup services container **301**, e.g., requesting an inventory of containerized applications **320**, receiving said inventory and storing it to management database **446**, generating storage management preferences for one or more of the inventoried applications (e.g., schedule, policies, retention, etc.), instructing backup services container **301** to track activity at one or more targeted applications **320**, transmitting trigger thresholds to backup services container **301**, communicating with data agent(s) **442** in backup services container **301**, etc. Storage manager **440** executes on a computing device comprising one or more hardware processors and computer memory (e.g., virtualized, non-virtualized, cloud-based, non-cloud, etc.) without limitation. More details are given in other figures.

(263) Management database **446** is analogous to management database **146** and additionally comprises information and data structures relating to containerized applications **320** and/or backup services containers **301**. Management database **446** is a logical part of storage manager **440**, but is not necessarily physically part of the computing device that hosts storage manager **440**.

(264) Discovery logic **450**, is a functional component of backup services container **301**. Discovery logic **450** determines which containerized applications **320** are actually present in the Kubernetes pod **310**. In some embodiments, discovery logic **450** reports its findings, including discovered application inventories and application attributes, to storage manager **440** in the proprietary data storage management system **302**. Storage manager **440**, which is generally responsible for controlling storage operations including backups, stores the information reported by discovery logic **450** (e.g., to management database **446**). Storage manager **440** further employs the information for generating preferences that apply to the discovered containerized applications **320**, e.g., storage policies, backup schedules, retention policies, etc. In some embodiments, storage manager **440** also establishes activity tracking of a targeted containerized application **320** to help trigger backup operations based on activity by the application (e.g., more frequent backups for busy applications, more frequent backups for applications generating large amounts of data, etc.) In some embodiments, backup services container **301** comprises activity monitoring logic (e.g., as part of discovery logic **450** and/or as a separate functional component) that tracks targeted containerized applications **320** and reports to the external backup system **302/303**, e.g., using pre-defined thresholds that are used for triggering backups.

(265) When a backup operation is triggered (e.g., by proprietary data storage management system **302**, by a third-party backup system **303**, etc.), backup services container **301** receives notice, e.g., in the form of a trigger, an instruction, a message, etc. Discovery logic **450**, responsive to the notice of a pending backup, determines and/or confirms which containerized applications **320** are operating in the present Kubernetes pod **310**. In some embodiments, discovery logic **450** communicates with external backup systems **302/303** via distinct interface logic **470**. In other embodiments, interface logic is incorporated into discovery logic **450**.

(266) Examples of information collected by discovery logic **450** include without limitation: container **319** identifiers and attributes, containerized applications **320** and attributes, storage volumes **330** and attributes, labels (typically assigned by users/programmers) for the various applications **320** and/or volumes **330**, metadata associated with the various applications **320** and/or volumes **330**, system logs, information inside file system(s), device entries, kernel representation of hardware resources (e.g., CPU, main memory/RAM), etc. Discovery logic **450** is designed to be broadly inclusive in finding configuration data structures (e.g., container configuration files/definitions), interpreting their contents, and discovering assets present within other containers **319** configured within Kubernetes pod **310** (e.g., applications **320**, volumes **330**). Discovery logic **450** is further designed to analyze metadata of assets, e.g., application metadata, to determine a type or kind and a version of each application **320** discovered in the pod. For example, an application **320** might be named “my_beautiful_paintings,” which does not readily convey what kind of application it is. Discovery logic **450** examines metadata (e.g., configuration parameters associated with the application) that indicate that this is an instance of a MySQL DBMS image version 5.6. Thus, discovery logic **450** determines that the application's label is “my_beautiful_paintings”; the application's type is MySQL DBMS; and the application's version is 5.6. Discovery logic **450** will collect other attributes, such as a data path to where the database data is located, i.e., where DBMS “my_beautiful_paintings” writes to and reads from, and the size of the

database.

(267) Selection logic **460** is a functional component of backup services container **301**. Selection logic **460** selects and applies suitable backup preparation toolkits **401**. First, selection logic **460** determines which of the discovered containerized applications **320** require which of the pre-configured backup preparation toolkits **401**, if any. Some containerized applications **320** need not be expressly prepared for backup, and therefore will not require the services of a backup toolkit **401**. Next, selection logic **460** selects a backup toolkit **401** that is suitable for a target application **320**, e.g., toolkit **401-b** for MySQL DBMS **320-1**. Selection logic **460** establishes communications with the target containerized application **320**, e.g., via container **319** comprising application **320**, via direct communications with application **320**, etc. Selection **320** selects a suitable script/command and executes the script/command in the toolkit **401**, e.g., invoking certain commands via an API, which causes the target application **320** to be readied for backup. For example, a backup preparation script issues one or more API commands to quiesce a target DBMS or file system. In some embodiments, a file system is quiesced by commands issued to a file manager application that accesses/manages the file system. When selection logic **460** discovers that container **301** lacks a suitable backup toolkit **401** for a discovered application **320** it raises an alarm to the external backup system **302/303** and/or transmits an error message to storage manager **440**.

(268) Selection logic **460** reports to the external backup system **302/303** that the target application **320** is ready for backup (e.g., quiescent). Selection logic **460** waits for an indication that the backup operation has completed, e.g., receiving notice from external backup system **302/303** after a snapshot is taken. Selection logic **460** then releases application **320** from its backup-ready (e.g., quiescent) state to resume normal operations (e.g., applies an unquiesce command). Any number of target containerized applications **320** can be prepared for backup in this way using backup toolkits **401** in backup services container **301**. With a wide range of toolkits **401** pre-configured and readily available, backup services container **301** is capable of preparing a wide range of target applications **320** for backup without necessitating an upgrade to backup services container **301** or to containers **319** comprising containerized applications **320**. In some embodiments, selection logic **460** communicates with external backup systems **302/303** via distinct interface logic **470**. In other embodiments, interface logic is incorporated into selection logic **460**.

(269) Backup interface logic **470** is a functional component of backup services container **301**. Backup interface logic **470** communicates with one or more of: backup system **303**, proprietary data storage management system **302** (e.g., storage manager **440**), containers **319**, containerized applications **320**, without limitation. In some embodiments, interface logic **470** operates as a distinct component and in other embodiments interface logic **470** is incorporated into other parts of backup services container **310**, e.g., discovery logic **450**, selection logic **460**, backup toolkits **401**, etc., without limitation.

(270) FIG. 4B is a block diagram depicting some salient details of backup services container **301** and data storage management system **302**, wherein one or more data agents **442** are deployed in backup services container **301**. FIG. 4B is similar to FIG. 4A, except that one or more data agents **442** are deployed within backup services container **301** rather than in backup proxy **406**. A virtual server data agent (VSA) is an example of a data agent **442**, which is generally used for backing up virtual machine data. This data agent is used for backing up containerized applications **320** running in a virtualized environment, such as Docker containers. A file system data agent (e.g., Windows, UNIX/Linux, Macintosh, etc.) is another example of a data agent **442**, which is generally used for file system backups. Accordingly, storage manager **440**, which controls storage management operations in system **302**, communicates with data agent **442** and media agent **444** to perform a backup or a restore operation. Data agent **442** is also in communication with media agent **444** for transmitting data thereto in a backup operation or for receiving data therefrom in a restore operation. The communicative pathways among storage manager **440**, data agent **442**, and media agent **444** are depicted by the bi-directional arrows **114**, which are described in more detail elsewhere herein.

(271) FIG. 4C is a block diagram depicting some salient details of system **300** in which a backup proxy is configured and deployed as a container-orchestration (Kubernetes) pod **486** within a Kubernetes node **413**. FIG. 4C is similar to FIG. 3D, except that each of the other pods in the node

here comprises a backup services container **301** (thereby being defined as a pod **310**), and Kubernetes node **413** further comprises a specially configured “backup services pod” **486**, which acts as a backup proxy of proprietary data storage management system **302** and is in communication with storage manager **440**. Thus, Kubernetes node **413** is similar to node **313** and additionally comprises components such as pod **486** and containers **310**.

(272) In this embodiment, the components ordinarily configured in backup proxy **406** (see, e.g., FIG. 4A) are containerized and configured into a Kubernetes pod **486** (the “backup services pod”) that is part of the same Kubernetes node **413** as pods **310** comprising containerized applications **320** and backup services containers **301**. For example, data agents **442** and media agents **444** are configured in backup services pod **486**. In some embodiments, data storage for secondary copies **116** also is part of backup services pod **486**, but the invention is not so limited. In some embodiments, storage manager **440** also is configured within backup services pod **486**, thus deploying proprietary data storage management system **302** within a Kubernetes node **413**. In all these various embodiments of backup services pod **486**, the proprietary features of system **302** are made portable under the Kubernetes orchestration umbrella of node **413**.

(273) FIG. 4D is a block diagram depicting some details of a backup services pod **486** deployed in a Kubernetes node **413**. Illustratively, pod **486** comprises any number of data agents **442** (e.g., VSA, Windows file system, UNIX/Linux file system, Macintosh file system, etc.); any number of media agents **444**; and any number of data storage volumes **330**. There is no limitation on how many and what kinds of data agents **442** are configured in any given backup services pod **486**. Likewise, there is no limitation in how many media agents **444** and volumes **330** are configured in any backup services pod **486**.

(274) In some embodiments (not shown here), storage manager **440** and management database **446** are also configured within backup services pod **486**, without limitation, thus deploying proprietary data storage management system **302** within a Kubernetes node **413**.

(275) FIG. 5 is a block diagram depicting some salient details of and logical information paths to/from backup services container **301** and storage manager **440**. The present figure depicts: backup services container **301** comprising backup toolkits **401**, discovery logic **450**, selection logic **460**, and interface logic **470**; storage manager **440** comprising backup preparation logic **541**, administration logic **543**, management database **446**, and container interface logic **545**; execution resources **501**; and backup preparation scripts **502**. The dotted bi-directional arrows indicate logical communication pathways between depicted components to enhance the reader's understanding of the present disclosure. The dotted unidirectional arrows indicate logical pathways connecting sources **501/502** with backup toolkits **401**.

(276) Execution resources **501** include any execution environment utilities needed to enable scripts **502** to run. Execution resources **501** are also referred to herein as “enabling utilities”. Example resources **501** include the Python programming language and its standard library. Python and python resources are well known in the art. Another example of resources **501** is the C standard library for the C programming language, sometimes referred to colloquially as “C Runtime” or “Runtime C.” Runtime C is well known in the art. There is no limitation on execution resources **501** and sources for such enabling utilities, whether the sources are proprietary, open source, or any combination thereof. The dotted unidirectional arrows from resources **501** to one or more backup toolkits **401** depict a logical flow of resources (enabling utilities) into toolkits **401**. It will be up to the implementers of an embodiment to equip each backup toolkit **401** with appropriate execution resources **501** sufficient to enable execution of the scripts **502** needed to run pre-backup and post-backup. For example, backup preparation scripts written in the C programming language would require Runtime C within toolkit **401** in order to guarantee that the C scripts will run as needed. On the other hand, backup preparation scripts written in python would require the python execution environment within toolkit **401** in order to guarantee that the python scripts will run as needed. And so on in regard to other programming languages and corresponding execution resources (enabling utilities) therefor.

(277) Backup preparation scripts **502** are programs that cause certain target applications to become quiescent and/or to unquiesce. Typically, a script **502** will access an application **320**, for example by

using an API, and the script **502** will issue one or more commands to the application **320** to make it quiesce. The same or a different script **502** within toolkit **401** will issue one or more commands to the application **320** to unquiesce it and resume operations after the backup operation completes. Implementation of script or scripts **502** for a given target application **320** will vary from one application to another and will be up to the implementers of an embodiment to ensure that scripts **502** are compatible with the target application's API and operational for the quiescing and unquiescing functions. There is no limitation on backup preparation scripts **502** and sources of such scripts, whether proprietary, open source, or any combination thereof. The dotted unidirectional arrows from scripts **502** to one or more backup toolkits **401** depict a logical flow of scripts into toolkits **401**. It will be up to the implementers of an embodiment to equip each backup toolkit **401** with appropriate scripts **502** accompanied by appropriate execution resources (enabling utilities) **501**.

(278) Some scripts **502** are not directed to quiescing and unquiescing and are instead used for discovery functions (see discovery logic **450**). Such discovery functions collect information about containerized assets (e.g., applications **320**, volumes **330**) and/or extract information from those assets (e.g., logs, configuration files, transaction history, etc.). These scripts also use the targeted application's API as appropriate and are also part of toolkit **401**.

(279) Backup preparation logic **541** is a functional component of storage manager **440**. Backup preparation logic **541** comprises routines for processing information obtained from discovery logic **450**. For example, logic **541** catalogues application inventories and attributes into a master catalog (not shown) that indexes all containerized applications **320**, their containers **319**, and respective Kubernetes pods **310**. The catalog comprises a list of applications **320**, whether they need preparation for backup or not. For example, logic **541** devises a backup schedule for the inventoried applications **320**, based on one or more backup criteria, e.g., time of day/week, amount of data to be backed up, type of application, etc. In some embodiments, logic **541** causes the criteria for backup to be stored in schedule preferences in management database **446**. In some embodiments, thresholds for triggering backups of certain applications **320** are transmitted by backup preparation logic **541** to discovery logic **450** for monitoring activity at the target application **320**, e.g., determining when a certain amount of data storage has been reached, determining when a certain amount of data has changed from a previous backup, etc. In some embodiments the catalog is stored in management database **446**.

(280) Administration logic **543** is a functional component of storage manager **440**. Backup preparation logic **541** illustratively implements administrative changes via administrative logic **543**, which stores catalogs, preferences, etc., to management database **446**. For example, for each discovered containerized application **320**, backup preparation logic **541** causes administrative logic **543** to create a "subclient" entity on management database **446**. Once the subclient entity is created, preferences are further created for it, e.g., backup schedules/triggers, media agents to use for backup, destination storage for secondary copies **116**, retention preferences for the secondary copies **116**, etc. These administrative actions are illustratively orchestrated by backup preparation logic **541** and implemented by administrative logic **543**, though the invention is not so limited.

(281) Container interface logic **545** is a functional component of storage manager **440** that is responsible for establishing and maintaining communications with container-based interface logic **470**. Notices of pending backup operations are transmitted by storage manager to container **301** using interface logic **545** and received by interface logic **470** at the container. Conversely, communications from backup services container **301** are transmitted by interface logic **470** and received by interface logic **545** at storage manager **440**. In alternative embodiments, communications between storage manager **440** and container **301** are maintained through other modules and means, without limitation.

(282) Functional logic modules **541**, **543**, and **545** are depicted as distinct modules here to enhance the reader's understanding of the present disclosure but the invention is not limited to having distinct functional modules and the features described herein can be implemented within one or more other parts of storage manager **440**, without limitation.

(283) FIG. **6** depicts some salient operations of a method **600** according to an illustrative embodiment.

(284) At block **602**, a backup services container **301** is configured, e.g., by storage manager **440**. Each backup services container **301** comprises wide-ranging content, including backup preparation toolkits

401 (e.g., scripts **502** comprising commands and enabling execution resources **501**). See also FIGS. **4A**, **4B**, **5**.

(285) At block **604**, backup services container **301** is installed in container-orchestration pod (e.g., Kubernetes pod) **310** comprising one or more containers **319** comprising containerized applications **320**. The installation is performed by one or more of storage manager **440** (e.g., using backup preparation logic **541**), by other programming tools, by manual intervention, etc., without limitation. Notably, installing backup services container **301** into a pod **310** does not affect the content or configurations of other containers in the pod (e.g., **319**).

(286) At block **606**, method **600** forks depending on whether backup services container **301** interoperates with a third-party backup system **303** (control passes to block **610**) or with proprietary data storage management system **302** (control passes to block **608**). Thus, depending on the nature of the system that initiates backups, certain additional features (e.g., proprietary services in block **608**) are provided.

(287) At block **608**, backup services container **301** in conjunction with proprietary data storage management system **302** performs certain enhanced services, e.g., discovery operations and reporting. More details are given in another figure. This block is skipped when a third-party backup system **303** is the backup initiator.

(288) At block **610**, on receiving notice of a pending backup operation, backup services container **301** prepares for backup selected containerized application(s) **320** in pod **310**. More details are given in another figure.

(289) At block **612**, the system that initiated the backup operation (e.g., **302**, **303**) executes the backup operation. In the present context and for purposes of the present disclosure, the backup operation is typically a snapshot taken of the targeted application's data, though the invention is not so limited. Taking of the snapshot, whether hardware snapshot or software snapshot, is generally very fast. Once the snapshot has been taken, backup services container **301** is notified to that effect (e.g., by storage manager **440**, by third-party system **303**, etc.). Storage manager **440** is specially configured to communicate with backup services container **301** (e.g., using container interface **545**, backup preparation logic **541**, and/or administrative logic **543**, without limitation). Subsequent backup processing of the snapshot in order to generate secondary copies **116** does not require the targeted containerized applications **320** to be in a quiescent state. In some embodiments, the backup operation is not snapshot-based and backup services container **301** is notified of backup completion at a suitable time.

(290) At block **614**, backup services container **301**, having received notice of the completion of the backup operation, releases the selected containerized application(s) **320**, after which control passes back to block **610**. Releasing the application(s) from their backup-ready state comprises causing application(s) **320** to resume ordinary operations, e.g., by sending an unquiesce command using selected toolkit **401**. Accordingly, selection logic **460** causes the selected backup toolkit **401** to run a post-backup script or scripts **502** that will cause application **320** to resume normal operations (i.e., to be unquiesced). To execute successfully, script or scripts **502** use the execution resources **501** provided in toolkit **401**.

(291) Backup processing of the snapshot in order to generate secondary copies **116** is described in more detail elsewhere herein. In proprietary data storage management system **302**, as in system **100**, a data agent **442/142** and a media agent **444/144** operate under direction of storage manager **440** to use the snapshot as a data source, process it according to preferences (e.g., identifying changed data for an incremental backup, compressing, encrypting, deduplicating, etc.), arrange the resulting data into a suitable secondary copy format, and store the secondary copy or copies **116** to prescribed storage resources, e.g., **108**, as well as performing any indexing operations.

(292) In embodiments wherein data agents **142/442** and media agents **144/444** operate in backup proxy **406** (see, e.g., FIG. **4A**), the snapshot taken during the backup operation is preferably stored (not shown) at the backup proxy **406** in order to optimize snapshot processing performance for backup, though the invention is not so limited. Illustratively, data agent **142/442** transmits data processed from the snapshot to media agent **444**, which ultimately generates the secondary copy **116**, indexes it, and

stores it to one or more storage resources (e.g., on-premises storage, remote data center storage, cloud storage, etc.)

(293) In embodiments that include a data agent **442** within backup services container **301** (see, e.g., FIG. **4B**), storage manager **440** and media agent **444** continue to maintain communications with backup services container **301** throughout data agent's **442** being active in the making of a secondary copy **116**, as data agent **442** processes source data from a snapshot, which is illustratively captured (not shown) into backup services container **301**. Performance efficiencies are realized by storing the snapshot within the same container **301** as data agent **442** appointed to process the snapshot, though the invention is not so limited. Illustratively, data agent **442** transmits data processed from the snapshot to media agent **444**, which ultimately generates the secondary copy **116**, indexes it, and stores it to one or more storage resources (e.g., on-premises storage, remote data center storage, cloud storage, etc.)

(294) In embodiments that include a backup services pod **486** (see, e.g., FIG. **4C**, FIG. **4D**), the snapshot taken during the backup operation is illustratively stored in the backup services pod **486**, e.g., in a storage volume **330**, and processed therein by data agent **442** and media agent **444**. Performance efficiencies are realized by storing the snapshot within the same pod **486** as the data agent **442** appointed to process the snapshot, though the invention is not so limited. Illustratively, within backup services pod **486**, data agent **442** transmits data processed from the snapshot to media agent **444**, which ultimately generates the secondary copy **116**, indexes it, and stores it to one or more storage resources within pod **486** and/or outside pod **486**, without limitation.

(295) Method **600** operates with one or more arrangements of proprietary data storage management system **302**, including data agents, media agents, and storage manager, as shown in FIGS. **4A**, **4B**, **4C**, and/or **4D**. Method **600** also or alternatively operates with third-party backup systems **303** as shown in FIG. **3B**, without limitation.

(296) FIG. **7** depicts some salient details of block **608** in method **600**. Block **608** is generally directed to performing enhanced services made possible by interoperating with proprietary data storage management system **302**.

(297) At block **702**, discovery logic **450** is initiated (launched, started, triggered, etc.) within backup services pod **301**. Initiation of discovery logic **450** illustratively is started in response to a request (command, instruction, message, etc.) received from storage manager **440**. In other embodiments, discovery logic **450** launches itself upon establishing communications with proprietary data storage management system **302**, e.g., with storage manager **440**.

(298) At block **704**, which executes discovery logic **450**, information about contents of Kubernetes pod **310** is collected. Examples of information collected by discovery logic **450** includes, container **319** identifiers and attributes, containerized applications **320** and attributes, storage volumes **330** and attributes, labels (typically assigned by users/programmers) for the various applications **320** and/or volumes **330**, metadata associated with the various applications **320** and/or volumes **330**, system logs, information inside file system(s), device entries, kernel representation of hardware resources (e.g., CPU, main memory/RAM), etc. Discovery logic **450** is designed to be broadly inclusive in finding configuration data structures (e.g., container configuration files/definitions), interpreting their contents, and discovering assets present within other containers **319** configured within Kubernetes pod **310** (e.g., applications **320**, volumes **330**). Discovery logic **450** is further designed to analyze metadata of assets, e.g., application metadata, to determine a type or kind and a version of each application **320** discovered in the pod.

(299) At block **706**, discovery logic **450** reports discovery results to storage manager **440**, e.g., using interface logic **470**. There is no limitation on the form in which discovery logic **450** reports information to storage manager **440**. For example, in some embodiments, discovery logic **450** generates a catalog of discovered applications **320** accompanied by their respective attributes and/or metadata, and likewise generates a catalog of discovered storage volumes **330**, their respective attributes, and their associations with applications **320**.

(300) At block **708**, storage manager **440** stores discovery results to management database **446**. In some embodiments, storage manager **440** stores the catalogs as received from discovery logic **450**; in some embodiments, storage manager **440** pre-processes information received from discovery logic **450**

before storing it to management database **446**; and/or any combination thereof without limitation. (301) At block **710**, based on discovery results, storage manager **440** takes administrative action (e.g., using backup preparation logic **541**, administrative logic **543**, etc.). Accordingly, storage manager analyzes the discovered information obtained from discovery logic **450** and formulates one or more action plans relating to the discovered entities, e.g., generates clients corresponding to containers **319**; generates subclients corresponding to applications **320** and/or volumes **330**; generates backup and/or retention preferences for each subclient (e.g., appropriate data agent, appropriate media agent; destination storage for secondary copies, retention duration for secondary copies, etc.); generates backup schedules for subclients (e.g., time schedules, dynamic triggers, etc.); etc. without limitation. Thus, storage manager **440** pro-actively processes information gathered by discovery logic **450** to ensure that data discovered in the Kubernetes pod **310** is protected in a streamlined fashion going forward.

(302) Thus, block **710**, and more generally block **608**, represent a way of integrating the data (including applications **320** generating the data) discovered in a Kubernetes pod **310** into the larger data protection umbrella of proprietary data storage management system **302**. Notably, as evidenced by the operations in the present figure, applications **320** discovered by discovery logic **450** need not be of a kind that requires pre-backup preparation (e.g., quiescing). Thus, any and all data-related assets that discovery logic **450** reports to storage manager **440** can be protected by proprietary data storage management system **302**. Thus, discovery logic **450** is leveraged here for broader data protection purposes than preparing certain kinds of applications for backup.

(303) FIG. **8** depicts some salient details of block **610** in method **600**. Block **610** is generally directed to preparing application(s) **320** in Kubernetes pod **310** for backup on receiving notice of a backup operation.

(304) At block **802**, backup services container **301** receives notice of an impending backup operation. Notice is received from proprietary data storage management system **302** (e.g., from storage manager **440**), or from a third-party backup system, e.g., **303**. Illustratively, notice is received by backup system interface logic **470**, which triggers discovery logic **450**, but the invention is not so limited.

(305) At block **804**, discovery logic **450** executes in backup services container **301**. Discovery logic **450** discovers contents (assets) of Kubernetes pod **310**, e.g., by analyzing container definitions and other configuration data structures. Discovery logic **450** identifies containerized applications **320** configured in pod **310**.

(306) At block **806**, discovery logic **450** illustratively reports its findings to selection logic **460**, though the invention is not so limited. See also FIG. **4A** and block **704** in FIG. **7**. Selection logic **460** is thus triggered to execute by reporting received from discovery logic **450**. Selection logic **460** analyzes each containerized application **320**. For each application **320**, selection logic **460** determines whether the application **320** needs to be prepared for backup (e.g., quiesced). On determining that a given containerized application is of a type that does not need pre-backup preparation, control passes to block **810**, because no further preparation is needed.

(307) On determining that a given containerized application is of a type that needs pre-backup preparation, selection logic **460** determines which of the backup toolkits **401** in backup services container **301** is suitable for performing the preparation. For example, selection logic **460** determines that an application **320** is of a type MySQL DBMS version 5.6 and finds a MySQL backup toolkit (e.g., **401-b**) that is suitable therefor. Selection logic **460** is specially configured to match the target application **320** to a suitable backup toolkit **401** that is suitable not just for the target application's type (e.g., MySQL) but also for its version (e.g., 5.6), because APIs change over time and can become outdated. Thus, having the intelligence to properly discover applications' attributes, and match applications **320** to backup toolkits **401** is one of the key features provided by backup services container **301**.

(308) At block **808**, selection logic **460** causes the selected backup toolkit **401** to establish communications with the target containerized application **320** and to run the pre-backup script or scripts **502** that will cause the application **320** to be ready for backup (e.g., quiesced). The script or scripts **502** will use the execution resources **501** provided in toolkit **401** to execute successfully.

(309) At block **810**, selection logic **460** (e.g., using interface logic **470**) reports to the initiating backup system (e.g., **302** using storage manager **440**, **303**) that the targeted application **320** is ready for the backup operation. Control passes to block **806** for processing another containerized application **320**.

(310) It should be noted here that information in the notice of backup operation received at block **802** may indicate one or more specific containers **319** or applications **320** or volumes **330** that are to be backed up, though that level of specificity is not always provided in every notice. When the notice provides no specific backup targets, backup services container **301** will proceed through all discovered containerized applications **320** found in pod **310**. When the notice indicates a specific container **319**, discovery logic **450** and selection logic **460** limit themselves to assets in that particular container in order to save time and unnecessary processing. When the backup notice indicates a specific target application **320** (e.g., based on a backup policy at storage manager **440** triggering a backup for a particular client/subclient), discovery logic **450** limits itself to the container **319** that comprises the targeted application **320** and selection logic **460** analyzes only the specified application **320** targeted for backup. Likewise in regard to specified containers **330**. Thus, the illustrative backup services container **301** flexibly addresses only those backup operations pending from the initiating backup system **302/303**. In alternative embodiments, backup services container **301** runs discovery logic **450** through all assets in pod **301** and likewise runs selection logic **460**, thus taking the opportunity to confirm and update any previously collected information.

(311) In regard to the figures described herein, other embodiments are possible within the scope of the present invention, such that the above-recited components, steps, blocks, operations, messages, requests, queries, and/or instructions are differently arranged, sequenced, sub-divided, organized, and/or combined. In some embodiments, a different component may initiate or execute a given operation.

EXAMPLE EMBODIMENTS

(312) Some example enumerated embodiments of the present invention are recited in this section in the form of computer-implemented methods, computer-based (or computer-enabled) systems, and non-transitory computer-readable media, without limitation.

(313) According to an illustrative embodiment, a computer-implemented method comprises: generating a first container that is based on an operating system-level virtualization service, wherein the first container comprises: (i) executable discovery logic, (ii) a plurality of executable scripts and enabling utilities for executing the scripts, wherein each script is configured to prepare for backup one or more corresponding applications, and (iii) executable selection logic; adding the first container to a container-orchestration pod that comprises one or more other containers comprising one or more containerized applications, wherein components of the container-orchestration pod, including the first container and the one or more other containers, run on a computing environment comprising at least one hardware processor and computer memory; by the discovery logic that executes in the first container, based on an indication that a backup operation has been triggered, identifying at least a first containerized application among the one or more containerized applications; by the selection logic that executes in the first container, determining a first executable script that is suitable for preparing the first containerized application for backup; by the selection logic, causing the first executable script to: (a) use enabling utilities to execute in the first container, (b) access the first containerized application, and (c) prepare the first containerized application for the backup operation; by the selection logic, indicating that the first containerized application is ready for the backup operation; and by the selection logic, on receiving an indication that the backup operation has completed, releasing the first containerized application from a backup-ready state.

(314) The above-recited embodiment wherein the adding of the first container does not change configurations of the one or more other containers. The above-recited embodiment wherein the discovery logic, the plurality of scripts, the enabling utilities, and the selection logic, have visibility to computing resources allocated to the first container and lack visibility to computing resources allocated outside the first container. The above-recited embodiment wherein the indication that the backup operation has been triggered indicates that the container-orchestration pod is to be backed up. The above-recited embodiment wherein the indication that the backup operation has been triggered

indicates that a second container in the container-orchestration pod is to be backed up. The above-recited embodiment wherein the indication that the backup operation has been triggered indicates that the first containerized application in the container-orchestration pod is to be backed up. The above-recited embodiment wherein the selection logic identifies the first containerized application as a target for the backup operation based on a backup indicator configured within the container-orchestration pod and associated with the first containerized application. The above-recited embodiment wherein to prepare the first containerized application for backup, the first executable script quiesces the first containerized application. The above-recited embodiment wherein the first containerized application is a database management system. The above-recited embodiment wherein the first containerized application is a file manager associated with a file system. The above-recited embodiment wherein the indication that the backup operation has completed is received after a snapshot is taken of data associated with the first containerized application. The above-recited embodiment wherein the container-orchestration pod is part of a node that executes as a service in a cloud computing account. The above-recited embodiment wherein the container-orchestration pod is part of a node that executes on a computing device comprising one or more hardware processors and computer memory. The above-recited embodiment wherein the container-orchestration pod is based on Kubernetes technology.

(315) The above-recited embodiment wherein the first container is a Docker container. The above-recited embodiment wherein the first container and the one or more other containers are Docker containers. The above-recited embodiment wherein the indication that the backup operation has been triggered is received from a data storage management system, and wherein the backup operation generates one or more secondary copies of data associated with the first containerized application. The above-recited embodiment wherein the indication that the backup operation has been triggered is received from a storage manager component of a data storage management system, wherein the storage manager controls storage operations including the backup operation, and wherein the backup operation generates one or more secondary copies of data associated with the first containerized application. The above-recited embodiment wherein the indication that the backup operation has been triggered is received from a backup system that operates outside the container-orchestration pod. The above-recited embodiment wherein the container-orchestration pod operates in a cloud computing environment. The above-recited embodiment wherein the container-orchestration pod operates in a computing environment that comprises one or more hardware processors and computer memory. The above-recited embodiment further comprising: by the discovery logic, collecting a plurality of attributes about the first containerized application in the container-orchestration pod; and transmitting the plurality of attributes about the first containerized application to a data storage management system that performs the backup operation. The above-recited embodiment further comprising: by the discovery logic, collecting a plurality of attributes about the first containerized application in the container-orchestration pod; transmitting the plurality of attributes about the first containerized application to a data storage management system that performs the backup operation; and based on the plurality of attributes, generating one or more preferences at the data storage management system for protecting data associated with the first containerized application. The above-recited embodiment further comprising: by the discovery logic, collecting information associated with a second container among the one or more other containers in the container-orchestration pod; and transmitting the information to a data storage management system that performs the backup operation. The above-recited embodiment wherein a storage manager that controls storage operations in the data storage management system, including the backup operation, uses the information collected by the discovery logic in the first container to administer preferences within the data storage management system for protecting data in the container-orchestration pod. The above-recited embodiment wherein a storage manager that controls storage operations in the data storage management system, including the backup operation, uses the information collected by the discovery logic, including an inventory of containerized applications, to administer preferences within the data storage management system for protecting data associated with the containerized applications in the inventory. The above-recited embodiment wherein a storage manager that controls storage operations in the data storage management system, including the backup operation, uses the information collected by the discovery

logic in the first container to administer entities within the data storage management system corresponding to data stored in the container-orchestration pod, and to control backup operations for the entities. The above-recited embodiment wherein a storage manager that controls storage operations in the data storage management system, including the backup operation, uses the information collected by the discovery logic in the first container to administer backup schedules within the data storage management system for protecting data in the container-orchestration pod. The above-recited embodiment wherein a storage manager that controls storage operations in the data storage management system, including the backup operation, uses the information collected by the discovery logic in the first container to perform live browse operations upon one or more secondary copies generated by the backup operation. The above-recited embodiment wherein a storage manager that controls storage operations in the data storage management system, including the backup operation, uses the information collected by the discovery logic in the first container, including information about data storage volumes attached to the first containerized application, to perform live browse operations upon one or more secondary copies of data stored in the data storage volumes that were generated by the backup operation.

(316) The above-recited embodiment further comprising: by the discovery logic, collecting information associated with a second container among the one or more other containers in the container-orchestration pod; and transmitting the information to a data storage management system that performs the backup operation, wherein the information comprises one or more of: information about applications in the second container, information about data storage in the second volume, metadata associated with the second container, system logs in the second container, information about kernel representations of hardware allocated to the second container.

(317) According to another embodiment a computer-based system comprises: a backup services container that is based on an operating system-level virtualization service, wherein the first container comprises: (i) executable discovery logic, (ii) a plurality of executable scripts and enabling utilities for executing the scripts, wherein each script is configured to prepare for backup one or more corresponding applications, and (iii) executable selection logic; wherein the backup services container is configured within a container-orchestration pod that comprises one or more other containers comprising one or more containerized applications; wherein components of the container-orchestration pod, including the backup services container and the one or more other containers, run on a computing environment comprising at least one hardware processor and computer memory; wherein the backup services container is configured to: based on an indication that a backup operation has been triggered, use the discovery logic to identify at least a first containerized application among the one or more containerized applications, use the selection logic to determine a first executable script that is suitable for preparing the first containerized application for backup, use the selection logic to cause the first executable script to: (a) use enabling utilities to execute in the backup services container, (b) access the first containerized application, and (c) prepare the first containerized application for the backup operation, use the selection logic to indicate that the first containerized application is ready for the backup operation, and on receiving an indication that the backup operation has completed, use the selection logic to release the first containerized application from a backup-ready state.

(318) The above-recited embodiment further comprising a data storage management system comprising a storage manager that controls storage operations in the data storage management system, including the backup operation; and wherein the storage manager executes on one of: a computing device comprising one or more hardware processors, a virtual machine executing on a computing device comprising one or more hardware processors. The above-recited embodiment wherein the backup services container is further configured to: use the discovery logic to collect information associated with a second container among the one or more other containers in the container-orchestration pod; and transmit the information to the data storage management system that performs the backup operation. The above-recited embodiment wherein the storage manager uses the information collected by the discovery logic to administer preferences within the data storage management system for protecting data in the container-orchestration pod. The above-recited embodiment wherein the storage manager uses the information collected by the discovery logic,

including an inventory of containerized applications, to administer preferences within the data storage management system for protecting data associated with the containerized applications in the inventory. The above-recited embodiment wherein the storage manager uses the information collected by the discovery logic in the first container to administer entities within the data storage management system corresponding to data stored in the container-orchestration pod, and to control backup operations for the entities. The above-recited embodiment wherein the storage manager uses the information collected by the discovery logic in the first container to administer backup schedules within the data storage management system for protecting data in the container-orchestration pod. The above-recited embodiment wherein the storage manager uses the information collected by the discovery logic in the first container to perform live browse operations upon one or more secondary copies generated by the backup operation. The above-recited embodiment wherein the storage manager uses the information collected by the discovery logic in the first container, including information about data storage volumes attached to the first containerized application, to perform live browse operations upon one or more secondary copies of data stored in the data storage volumes that were generated by the backup operation. The above-recited embodiment wherein the backup services container is further configured to use the discovery logic to collect information associated with a second container among the one or more other containers in the container-orchestration pod; and wherein the information comprises one or more of: information about applications in the second container, information about data storage in the second volume, metadata associated with the second container, system logs in the second container, information about kernel representations of hardware allocated to the second container.

(319) According to yet another exemplary embodiment, a method comprises: generating a first container that is based on an operating system-level virtualization service, wherein the first container comprises: (i) executable discovery logic, (ii) a plurality of executable scripts and enabling utilities for executing the scripts, wherein each script is configured to prepare for backup one or more corresponding applications, and (iii) executable selection logic; adding the first container to a first container-orchestration pod that comprises one or more other containers comprising one or more containerized applications, wherein components of the container-orchestration pod, including the first container and the one or more other containers, run on a computing environment comprising at least one hardware processor and computer memory, and wherein the adding of the first container does not change configurations of the one or more other containers; generating a second container-orchestration pod (backup services pod) comprising one or more data agents and one or more media agents for generating secondary copies of data from the one or more containerized applications; by the discovery logic that executes in the first container, based on an indication that a backup operation has been triggered, identifying at least a first containerized application among the one or more containerized applications; by the selection logic that executes in the first container, determining a first executable script that is suitable for preparing the first containerized application for backup; by the selection logic, causing the first executable script to: (a) use enabling utilities to execute in the first container, (b) access the first containerized application, and (c) prepare the first containerized application for the backup operation; by the selection logic, indicating that the first containerized application is ready for the backup operation; and by the selection logic, on receiving an indication that the backup operation has completed, releasing the first containerized application from a backup-ready state; and by one of the data agents and one of the media agents, generating a secondary copy of data generated by the first containerized application.

(320) The above-recited embodiment wherein a data storage management system comprises a storage manager that controls storage operations in the data storage management system, including the backup operation; and wherein the storage manager executes on one of: the second container-orchestration pod (backup services pod), a computing device comprising one or more hardware processors, and a virtual machine executing on a computing device comprising one or more hardware processors. The above-recited embodiment wherein the backup services container is further configured to: use the discovery logic to collect information associated with a second container among the one or more other containers in the container-orchestration pod; and transmit the information to the data storage management system that performs the backup operation. The above-recited embodiment wherein the storage

manager uses the information collected by the discovery logic to administer preferences within the data storage management system for protecting data in the container-orchestration pod. The above-recited embodiment wherein the storage manager uses the information collected by the discovery logic, including an inventory of containerized applications, to administer preferences within the data storage management system for protecting data associated with the containerized applications in the inventory. The above-recited embodiment wherein the storage manager uses the information collected by the discovery logic in the first container to administer entities within the data storage management system corresponding to data stored in the container-orchestration pod, and to control backup operations for the entities. The above-recited embodiment wherein the storage manager uses the information collected by the discovery logic in the first container to administer backup schedules within the data storage management system for protecting data in the container-orchestration pod. The above-recited embodiment wherein the storage manager uses the information collected by the discovery logic in the first container to perform live browse operations upon one or more secondary copies generated by the backup operation. The above-recited embodiment wherein the storage manager uses the information collected by the discovery logic in the first container, including information about data storage volumes attached to the first containerized application, to perform live browse operations upon one or more secondary copies of data stored in the data storage volumes that were generated by the backup operation. The above-recited embodiment wherein the backup services container is further configured to use the discovery logic to collect information associated with a second container among the one or more other containers in the container-orchestration pod; and wherein the information comprises one or more of: information about applications in the second container, information about data storage in the second volume, metadata associated with the second container, system logs in the second container, information about kernel representations of hardware allocated to the second container.

(321) According to another illustrative embodiment, a method comprises: generating a first container that is based on an operating system-level virtualization service, wherein the first container comprises: (i) executable discovery logic, (ii) a plurality of executable scripts and enabling utilities for executing the scripts, wherein each script is configured to prepare for backup one or more corresponding applications, and (iii) executable selection logic; adding the first container to a first container-orchestration pod that comprises one or more other containers comprising one or more containerized applications, wherein components of the container-orchestration pod, including the first container and the one or more other containers, run on a computing environment comprising at least one hardware processor and computer memory; generating a second container-orchestration pod (backup services pod) comprising components of a data storage management system for generating secondary copies of data generated by the one or more containerized applications, wherein the components include: a storage manager, one or more data agents, and one or more media agents; by the discovery logic that executes in the first container, based on an indication from the storage manager that a backup operation has been triggered, identifying at least a first containerized application among the one or more containerized applications; by the selection logic that executes in the first container, determining a first executable script that is suitable for preparing the first containerized application for backup; by the selection logic, causing the first executable script to: (a) use enabling utilities to execute in the first container, (b) access the first containerized application, and (c) prepare the first containerized application for the backup operation; by the selection logic, indicating that the first containerized application is ready for the backup operation; and by the selection logic, on receiving an indication that the backup operation has completed, releasing the first containerized application from a backup-ready state; and by the storage manager, instructing one of the data agents and one of the media agents to generate a secondary copy of data generated by the first containerized application, wherein the secondary copy of data is based on a snapshot taken by the backup operation.

(322) The above-recited embodiment wherein the adding of the first container does not change configurations of the one or more other containers. The above-recited embodiment wherein the one of the media agents stores the secondary copy to a data storage volume in the second container-orchestration pod (backup services pod). The above-recited embodiment wherein the one of the media

agents stores the secondary copy to a data storage volume outside the second container-orchestration pod (backup services pod). The above-recited embodiment wherein at least some of the information collected by the discovery logic in the first container is used in a live browse operation of the secondary copy. The above-recited embodiment wherein at least some of the information collected by the discovery logic in the first container is used for viewing a data file generated by the first containerized application in a live browse operation of the secondary copy. The above-recited embodiment wherein at least some of the information collected by the discovery logic in the first container is used for mounting a data storage volume in a live browse operation of the secondary copy. The above-recited embodiment wherein at least some of the information collected by the discovery logic in the first container is used for content indexing of the secondary copy.

(323) In other embodiments, a system or systems operates according to one or more of the methods and/or computer-readable media recited in the preceding paragraphs. In yet other embodiments, a method or methods operates according to one or more of the systems and/or computer-readable media recited in the preceding paragraphs. In yet more embodiments, a non-transitory computer-readable medium or media causes one or more computing devices having one or more processors and computer-readable memory to operate according to one or more of the systems and/or methods recited in the preceding paragraphs.

Terminology

(324) Conditional language, such as, among others, “can,” “could,” “might,” or “may,” unless specifically stated otherwise, or otherwise understood within the context as used, is generally intended to convey that certain embodiments include, while other embodiments do not include, certain features, elements and/or steps. Thus, such conditional language is not generally intended to imply that features, elements and/or steps are in any way required for one or more embodiments or that one or more embodiments necessarily include logic for deciding, with or without user input or prompting, whether these features, elements and/or steps are included or are to be performed in any particular embodiment.

(325) Unless the context clearly requires otherwise, throughout the description and the claims, the words “comprise,” “comprising,” and the like are to be construed in an inclusive sense, as opposed to an exclusive or exhaustive sense, i.e., in the sense of “including, but not limited to.” As used herein, the terms “connected,” “coupled,” or any variant thereof means any connection or coupling, either direct or indirect, between two or more elements; the coupling or connection between the elements can be physical, logical, or a combination thereof. Additionally, the words “herein,” “above,” “below,” and words of similar import, when used in this application, refer to this application as a whole and not to any particular portions of this application. Where the context permits, words using the singular or plural number may also include the plural or singular number respectively. The word “or” in reference to a list of two or more items, covers all of the following interpretations of the word: any one of the items in the list, all of the items in the list, and any combination of the items in the list. Likewise the term “and/or” in reference to a list of two or more items, covers all of the following interpretations of the word: any one of the items in the list, all of the items in the list, and any combination of the items in the list.

(326) In some embodiments, certain operations, acts, events, or functions of any of the algorithms described herein can be performed in a different sequence, can be added, merged, or left out altogether (e.g., not all are necessary for the practice of the algorithms). In certain embodiments, operations, acts, functions, or events can be performed concurrently, e.g., through multi-threaded processing, interrupt processing, or multiple processors or processor cores or on other parallel architectures, rather than sequentially.

(327) Systems and modules described herein may comprise software, firmware, hardware, or any combination(s) of software, firmware, or hardware suitable for the purposes described. Software and other modules may reside and execute on servers, workstations, personal computers, computerized tablets, PDAs, and other computing devices suitable for the purposes described herein. Software and other modules may be accessible via local computer memory, via a network, via a browser, or via other means suitable for the purposes described herein. Data structures described herein may comprise computer files, variables, programming arrays, programming structures, or any electronic information

storage schemes or methods, or any combinations thereof, suitable for the purposes described herein. User interface elements described herein may comprise elements from graphical user interfaces, interactive voice response, command line interfaces, and other suitable interfaces.

(328) Further, processing of the various components of the illustrated systems can be distributed across multiple machines, networks, and other computing resources. Two or more components of a system can be combined into fewer components. Various components of the illustrated systems can be implemented in one or more virtual machines, rather than in dedicated computer hardware systems and/or computing devices. Likewise, the data repositories shown can represent physical and/or logical data storage, including, e.g., storage area networks or other distributed storage systems. Moreover, in some embodiments the connections between the components shown represent possible paths of data flow, rather than actual connections between hardware. While some examples of possible connections are shown, any of the subset of the components shown can communicate with any other subset of components in various implementations.

(329) Embodiments are also described above with reference to flow chart illustrations and/or block diagrams of methods, apparatus (systems) and computer program products. Each block of the flow chart illustrations and/or block diagrams, and combinations of blocks in the flow chart illustrations and/or block diagrams, may be implemented by computer program instructions. Such instructions may be provided to a processor of a general purpose computer, special purpose computer, specially-equipped computer (e.g., comprising a high-performance database server, a graphics subsystem, etc.) or other programmable data processing apparatus to produce a machine, such that the instructions, which execute via the processor(s) of the computer or other programmable data processing apparatus, create means for implementing the acts specified in the flow chart and/or block diagram block or blocks. These computer program instructions may also be stored in a non-transitory computer-readable memory that can direct a computer or other programmable data processing apparatus to operate in a particular manner, such that the instructions stored in the computer-readable memory produce an article of manufacture including instruction means which implement the acts specified in the flow chart and/or block diagram block or blocks. The computer program instructions may also be loaded to a computing device or other programmable data processing apparatus to cause operations to be performed on the computing device or other programmable apparatus to produce a computer implemented process such that the instructions which execute on the computing device or other programmable apparatus provide steps for implementing the acts specified in the flow chart and/or block diagram block or blocks.

(330) Any patents and applications and other references noted above, including any that may be listed in accompanying filing papers, are incorporated herein by reference. Aspects of the invention can be modified, if necessary, to employ the systems, functions, and concepts of the various references described above to provide yet further implementations of the invention. These and other changes can be made to the invention in light of the above Detailed Description. While the above description describes certain examples of the invention, and describes the best mode contemplated, no matter how detailed the above appears in text, the invention can be practiced in many ways. Details of the system may vary considerably in its specific implementation, while still being encompassed by the invention disclosed herein. As noted above, particular terminology used when describing certain features or aspects of the invention should not be taken to imply that the terminology is being redefined herein to be restricted to any specific characteristics, features, or aspects of the invention with which that terminology is associated. In general, the terms used in the following claims should not be construed to limit the invention to the specific examples disclosed in the specification, unless the above Detailed Description section explicitly defines such terms. Accordingly, the actual scope of the invention encompasses not only the disclosed examples, but also all equivalent ways of practicing or implementing the invention under the claims.

(331) To reduce the number of claims, certain aspects of the invention are presented below in certain claim forms, but the applicant contemplates other aspects of the invention in any number of claim forms. For example, while only one aspect of the invention is recited as a means-plus-function claim under 35 U.S.C. sec. 112(f) (AIA), other aspects may likewise be embodied as a means-plus-function

claim, or in other forms, such as being embodied in a computer-readable medium. Any claims intended to be treated under 35 U.S.C. § 112(f) will begin with the words “means for,” but use of the term “for” in any other context is not intended to invoke treatment under 35 U.S.C. § 112(f). Accordingly, the applicant reserves the right to pursue additional claims after filing this application, in either this application or in a continuing application.

Claims

1. A computer-implemented method comprising: by a first container-orchestration pod that is configured to perform secondary copy operations of a containerized application, wherein the containerized application executes in a second container-orchestration pod, and wherein the containerized application comprises data stored in the second container-orchestration pod, performing operations comprising: causing a secondary copy operation to be executed on the data of the containerized application, thereby generating a secondary copy within the first container-orchestration pod, storing data of the secondary copy in a data storage volume that is configured at the first container-orchestration pod, and wherein the secondary copy is configured in a backup format that is distinct from a primary data format of the data of the containerized application; wherein the first container-orchestration pod and the second container-orchestration pod are configured in a computing environment that comprises one or more computer hardware processors.
2. The computer-implemented method of claim 1 further comprising: launching discovery logic at the second container-orchestration pod, based on communications with a storage manager that manages the secondary copy operation; by the discovery logic, identifying the containerized application at the second container-orchestration pod, and further identifying the data of the containerized application; and reporting to the storage manager, by the discovery logic, information that identifies the containerized application and the data of the containerized application, and further reporting attributes of the containerized application and of the data of the containerized application; wherein the storage manager executes on one or more computer hardware processors, and wherein the storage manager is configured to manage secondary copy operations of a data storage management system.
3. The computer-implemented method of claim 2 wherein the storage manager causes the launching of the discovery logic.
4. The computer-implemented method of claim 2 further comprising: generating data protection preferences for the containerized application and for the data of the containerized application, wherein the data protection preferences are based on the information and the attributes that are reported by and received from the discovery logic; wherein the secondary copy operation is based on the data protection preferences.
5. The computer-implemented method of claim 4 wherein the data protection preferences comprise one or more of: a data agent, a media agent, a destination storage for the secondary copy, and a retention duration for the secondary copy.
6. The computer-implemented method of claim 2, further comprising: integrating data protection for the containerized application and for the data of the containerized application into the data storage management system that is managed by the storage manager.
7. The computer-implemented method of claim 1, wherein processing the data of the secondary copy is performed by one or more of: a data agent configured in the first container-orchestration pod, and a media agent configured in the first container-orchestration pod; and wherein the data agent and the media agent are components of a data storage management system, which is configured to provide data protection for the containerized application and for the data of the containerized application.
8. The computer-implemented method of claim 1 wherein the first container-orchestration pod is configured as a backup services pod that comprises components of a data storage management system, wherein the data storage management system is configured to provide data protection for containerized applications of and for data within the second container-orchestration pod, including data protection for the containerized application and for the data of the containerized application.

9. The computer-implemented method of claim 1 further comprising: storing the secondary copy at one or more of: a first data storage volume within the first container-orchestration pod, and a second data storage volume outside of the first container-orchestration pod.
10. The computer-implemented method of claim 1, wherein first container-orchestration pod and the second container-orchestration pod are based on Kubernetes technology.
11. A computer-implemented system, the computer-implemented system comprising: one or more computing devices comprising computer hardware with one or more processors, the one or more computing devices configured to: by a first container-orchestration pod that is configured to perform secondary copy operations of a containerized application, wherein the containerized application executes in a second container-orchestration pod, and wherein the containerized application comprises data stored in the second container-orchestration pod, performing operations comprising: causing a secondary copy operation to be executed on the data of the containerized application, thereby generating a secondary copy within the first container-orchestration pod, storing data of the secondary copy in a data storage volume that is configured at the first container-orchestration pod, and wherein the secondary copy is configured in a backup format that is distinct from a primary data format of the data of the containerized application; wherein the first container-orchestration pod and the second container-orchestration pod are configured in a computing environment that comprises one or more computer hardware processors.
12. The computer-implemented system of claim 11, wherein the one or more computing devices configured to: launch discovery logic at the second container-orchestration pod, based on communications with a storage manager that manages the secondary copy operation; by the discovery logic, identify the containerized application at the second container-orchestration pod, and further identifying the data of the containerized application; and report to the storage manager, by the discovery logic, information that identifies the containerized application and the data of the containerized application, and further reporting attributes of the containerized application and of the data of the containerized application; wherein the storage manager executes on one or more computer hardware processors, and wherein the storage manager is configured to manage secondary copy operations of a data storage management system.
13. The computer-implemented system of claim 12, wherein the storage manager causes the launch of the discovery logic.
14. The computer-implemented system of claim 12, wherein the one or more computing devices configured to: generate data protection preferences for the containerized application and for the data of the containerized application, wherein the data protection preferences are based on the information and the attributes that are reported by and received from the discovery logic; wherein the secondary copy operation is based on the data protection preferences.
15. The computer-implemented system of claim 14, wherein the data protection preferences comprise one or more of: a data agent, a media agent, a destination storage for the secondary copy, and a retention duration for the secondary copy.
16. The computer-implemented system of claim 12, wherein the one or more computing devices configured to: integrate data protection for the containerized application and for the data of the containerized application into the data storage management system that is managed by the storage manager.
17. The computer-implemented system of claim 11, wherein processing the data of the secondary copy is performed by one or more of: a data agent configured in the first container-orchestration pod, and a media agent configured in the first container-orchestration pod; and wherein the data agent and the media agent are components of a data storage management system, which is configured to provide data protection for the containerized application and for the data of the containerized application.
18. The computer-implemented system of claim 11, wherein the first container-orchestration pod is configured as a backup services pod that comprises components of a data storage management system, wherein the data storage management system is configured to provide data protection for containerized applications of and for data within the second container-orchestration pod, including data protection for the containerized application and for the data of the containerized application.

19. The computer-implemented system of claim 11, wherein the one or more computing devices configured to: store the secondary copy at one or more of: a first data storage volume within the first container-orchestration pod, and a second data storage volume outside of the first container-orchestration pod.

20. The computer-implemented system of claim 11, wherein first container-orchestration pod and the second container-orchestration pod are based on Kubernetes technology.
