

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12399812
Kind Code	B2
Date of Patent	August 26, 2025
Inventor(s)	Wang; Shu et al.

Object sealing for cache coherency for shared memory

Abstract

This disclosure provides methods, devices, and systems for memory management. The present implementations more specifically relate to techniques for providing shared memory services over a Compute Express Link (CXL) fabric. In some aspects, a memory management system may include a shared memory (SM) manager and multiple SM libraries each associated with a respective host device coupled to a shared memory via a CXL fabric. In some implementations, an SM library may negotiate read or write “locks” with the SM manager for an object in shared memory, where a read-lock grants the associated host device read access to the object and a write-lock grants the associated host device write access to the object. In some implementations, the SM manager may grant any number of read-locks, but only one write-lock, to any object in shared memory so that the object becomes immutable once the host device relinquishes its write-lock for the object.

Inventors: Wang; Shu (Milpitas, CA), Beauchamp; Robert W. (Berlin, MA), Fan; Chenggong (Saratoga, CA)

Applicant: MemVerge, Inc. (Milpitas, CA)

Family ID: 1000008778286

Assignee: MemVerge, Inc. (Milpitas, CA)

Appl. No.: 18/358782

Filed: July 25, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240045804 A1	Feb. 08, 2024

Related U.S. Application Data

Publication Classification

Int. Cl.: **G06F12/02** (20060101); **G06F9/52** (20060101); **G06F9/54** (20060101); **G06F12/0815** (20160101)

U.S. Cl.:

CPC **G06F12/023** (20130101); **G06F9/526** (20130101); **G06F9/544** (20130101); **G06F12/0815** (20130101); G06F2213/0026 (20130101)

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6760815	12/2003	Traversat	711/E12.012	G06F 12/0276
7844973	12/2009	Dice	710/200	G06F 9/52
11042496	12/2020	BeSerra	N/A	G06F 13/4282
2003/0210248	12/2002	Wyatt	345/557	G06F 12/0888
2004/0255087	12/2003	Pudipeddi	711/147	G06F 12/0261
2012/0180074	12/2011	Krishnamurthy	719/315	G06F 21/14
2020/0019332	12/2019	Danilov	N/A	G06F 11/1076
2021/0157521	12/2020	Roy	N/A	G06F 3/067
2023/0236995	12/2022	Jayaraman	710/308	G06F 12/084
2024/0037026	12/2023	Price	N/A	G06F 3/0631
2024/0256348	12/2023	Pean	N/A	G06F 9/541

OTHER PUBLICATIONS

Norman Jouppi. "Cache Write Policies and Performance." May 1993. IEEE. SIGARCH Computer Architecture News. vol. 21. pp. 191-201. (Year: 1993). cited by examiner

Dr. Debendra Das Sharma and Ishwar Agarwal, "Compute Express Link™: The Breakthrough CPU-to-Device Interconnect CXL™," retrieved on Oct. 31, 2023 from <https://www.computeexpresslink.org/>, pp. 1-4. cited by applicant

Primary Examiner: Sadler; Nathan

Attorney, Agent or Firm: Paradice & Li LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION (1) This application claims priority and benefit under 35 USC § 119(e) to U.S. Provisional Patent Application No. 63/394,257, filed on Aug. 1, 2022, and to U.S. Provisional Patent Application No. 63/503,647, filed on May 22, 2023, which are incorporated herein by reference in their entireties.

TECHNICAL FIELD

(1) The present implementations relate generally to memory management in computer systems, and specifically to techniques for managing cache coherency for shared memory.

BACKGROUND OF RELATED ART

(2) Memory is generally used to store data for quick or immediate access by a computer system or processing resource. Many computer systems include a memory controller that reads data from, and writes data to, memory on behalf of various processing resources. The physical media that carries the data between the memory controller and the memory resources is generally referred to as a “memory bus.” Many existing memory buses (such as double data rate (DDR) memory buses) have physical constraints and timing protocols that require the memory resources to be located in relatively close proximity of the memory controller. As a result, many existing computer systems have local memory disposed on or within the computer system itself. Locally attached memory and processing resources are in the same “coherency domain.”

(3) New interconnect technologies are being developed to enable high-bandwidth, low-latency memory access over greater distances. For example, Compute Express Link (CXL) is an open standard interconnect that supports communications between a processor and memory over a Peripheral Component Interconnect Express (PCIe) link. As a result, CXL can extend coherency domains beyond small, localized fabrics to larger fabrics spanning multiple computer systems or remotely located devices. The CXL standard has led to improvements in memory usage and allocation (such as pooling of memory resources). As interconnect technologies continue to evolve, new memory management techniques are needed to support such advancements in memory usage and allocation.

SUMMARY

(4) This Summary is provided to introduce in a simplified form a selection of concepts that are further described below in the Detailed Description. This Summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to limit the scope of the claimed subject matter.

(5) One innovative aspect of the subject matter of this disclosure can be implemented in a method of memory management. The method includes creating an object in a memory shared by a plurality of host computing devices over a Compute Express Link (CXL) fabric; mapping a write-lock to the object in the shared memory responsive to creating the object, where the write-lock grants a first host computing device of the plurality of host computing devices exclusive access for writing to the object; receiving a request from the first host computing device to relinquish the write-lock; and sealing the object in the shared memory so that the object is immutable based on the request to relinquish the write-lock.

(6) Another innovative aspect of the subject matter of this disclosure can be implemented in a memory management system, including a processing system and a memory. The memory stores instructions that, when executed by the processing system, cause the memory management system to create an object in a memory shared by a plurality of host computing devices over a CXL fabric; map a write-lock to the object in the shared memory responsive to creating the object, where the write-lock grants a first host computing device of the plurality of host computing devices exclusive access for writing to the object; receive a request from the first host computing device to relinquish the write-lock; and seal the object in the shared memory so that the object is immutable based on the request to relinquish the write-lock.

(7) Another innovative aspect of the subject matter of this disclosure can be implemented in a

method of memory management. The method includes receiving a request to delete an object in a memory shared by a plurality of host computing devices; determining whether each host computing device of the plurality of host computing devices has ever accessed the object; flushing a respective cache associated with any host computing device of the plurality of host computing devices that has ever accessed the object; and deleting the object from the shared memory after flushing all the caches associated with any host computing device of the plurality of host computing devices that has ever accessed the object.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The present implementations are illustrated by way of example and are not intended to be limited by the figures of the accompanying drawings.
- (2) FIG. 1 shows a hardware diagram of an example computer system with pooled memory resources.
- (3) FIG. 2 shows a block diagram of an example computer system with pooled memory services, according to some implementations.
- (4) FIG. 3 shows a block diagram of an example memory management system, according to some implementations.
- (5) FIG. 4 shows a block diagram of an example host computing device, according to some implementations.
- (6) FIG. 5 shows another block diagram of an example memory management system, according to some implementations.
- (7) FIGS. 6A-6E show a series of operations that can be performed by a memory management system, according to some implementations.
- (8) FIG. 7 shows another block diagram of an example memory management system, according to some implementations.
- (9) FIG. 8 shows an illustrative flowchart depicting an example operation for memory management, according to some implementations.
- (10) FIG. 9 shows another illustrative flowchart depicting an example operation for memory management, according to some implementations.

DETAILED DESCRIPTION

(11) In the following description, numerous specific details are set forth such as examples of specific components, circuits, and processes to provide a thorough understanding of the present disclosure. The term “coupled” as used herein means connected directly to or connected through one or more intervening components or circuits. The terms “computer system,” “electronic system,” and “electronic device” may be used interchangeably to refer to any system capable of electronically processing information. Also, in the following description and for purposes of explanation, specific nomenclature is set forth to provide a thorough understanding of the aspects of the disclosure. However, it will be apparent to one skilled in the art that these specific details may not be required to practice the example embodiments. In other instances, well-known circuits and devices are shown in block diagram form to avoid obscuring the present disclosure. Some portions of the detailed descriptions which follow are presented in terms of procedures, logic blocks, processing and other symbolic representations of operations on data bits within a computer memory.

(12) These descriptions and representations are the means used by those skilled in the data processing arts to most effectively convey the substance of their work to others skilled in the art. In the present disclosure, a procedure, logic block, process, or the like, is conceived to be a self-consistent sequence of steps or instructions leading to a desired result. The steps are those requiring

physical manipulations of physical quantities. Usually, although not necessarily, these quantities take the form of electrical or magnetic signals capable of being stored, transferred, combined, compared, and otherwise manipulated in a computer system. It should be borne in mind, however, that all of these and similar terms are to be associated with the appropriate physical quantities and are merely convenient labels applied to these quantities.

(13) Unless specifically stated otherwise as apparent from the following discussions, it is appreciated that throughout the present application, discussions utilizing the terms such as “accessing,” “receiving,” “sending,” “using,” “selecting,” “determining,” “normalizing,” “multiplying,” “averaging,” “monitoring,” “comparing,” “applying,” “updating,” “measuring,” “deriving” or the like, refer to the actions and processes of a computer system, or similar electronic computing device, that manipulates and transforms data represented as physical (electronic) quantities within the computer system's registers and memories into other data similarly represented as physical quantities within the computer system memories or registers or other such information storage, transmission or display devices.

(14) In the figures, a single block may be described as performing a function or functions; however, in actual practice, the function or functions performed by that block may be performed in a single component or across multiple components, and/or may be performed using hardware, using software, or using a combination of hardware and software. To clearly illustrate this interchangeability of hardware and software, various illustrative components, blocks, modules, circuits, and steps have been described below generally in terms of their functionality. Whether such functionality is implemented as hardware or software depends upon the particular application and design constraints imposed on the overall system. Skilled artisans may implement the described functionality in varying ways for each particular application, but such implementation decisions should not be interpreted as causing a departure from the scope of the present disclosure. Also, the example input devices may include components other than those shown, including well-known components such as a processor, memory and the like.

(15) The techniques described herein may be implemented in hardware, software, firmware, or any combination thereof, unless specifically described as being implemented in a specific manner. Any features described as modules or components may also be implemented together in an integrated logic device or separately as discrete but interoperable logic devices. If implemented in software, the techniques may be realized at least in part by a non-transitory processor-readable storage medium including instructions that, when executed, performs one or more of the methods described above. The non-transitory processor-readable data storage medium may form part of a computer program product, which may include packaging materials.

(16) The non-transitory processor-readable storage medium may comprise random access memory (RAM) such as synchronous dynamic random-access memory (SDRAM), read only memory (ROM), non-volatile random access memory (NVRAM), electrically erasable programmable read-only memory (EEPROM), FLASH memory, other known storage media, and the like. The techniques additionally, or alternatively, may be realized at least in part by a processor-readable communication medium that carries or communicates code in the form of instructions or data structures and that can be accessed, read, and/or executed by a computer or other processor.

(17) The various illustrative logical blocks, modules, circuits and instructions described in connection with the embodiments disclosed herein may be executed by one or more processors (or a processing system). The term “processor,” as used herein may refer to any general-purpose processor, special-purpose processor, conventional processor, controller, microcontroller, and/or state machine capable of executing scripts or instructions of one or more software programs stored in memory.

(18) As described above, Compute Express Link (CXL) is an open standard interconnect that supports communications between a processor and memory over a Peripheral Component Interconnect Express (PCIe) link. A CXL system (or subsystem) includes a host computing device

(also referred to simply as a “host device”) and one or more CXL devices. The host device represents a root complex of the CXL system, whereas a CXL device can be any endpoint device coupled to the host device. Example CXL devices include Type 1 devices (such as accelerator devices without any locally-attached memory), Type 2 devices (such as accelerator devices with locally-attached memory that can be accessed by the host device), and Type 3 devices (such as passive memory expansion devices). A CXL fabric is a collection of one or more switches that rely on Hierarchy Based Routing (HBR) or Port Based Routing (PBR) messages to support flexible, low-latency switching between the host devices and CXL devices of a given CXL system.

(19) A coherency domain includes all memory and processing resources associated with a single coherent Host Physical Address (HPA) space. In existing computer systems, coherency domains were limited to small, localized fabrics built upon proprietary memory buses. The CXL standard defines a set of protocols that can be used to extend coherency domains to larger fabrics spanning multiple host devices, multiple CXL devices, or any combination thereof. As such, the CXL standard opens up new possibilities for storing and accessing data (such as memory pooling and sharing). Aspects of the present disclosure recognize that the speed and flexibility of the CXL architecture allows memory to be provided as a “service.” As used herein, the term “memory service” refers to any combination of hardware or software that manages or monitors the storage of data, and access thereto, on behalf of one or more host devices.

(20) For example, a CXL fabric can connect multiple host devices to the same memory resource(s). This allows data stored in such memory resources to be shared between the host devices. Many computing devices utilize local caches (which are much closer to the processor than other memory resources) to buffer data for quicker access. For example, when operating on shared data, a host device may load a copy of the data from shared memory to a local cache and subsequently read or modify the cached copy of the data. Thus, cache coherency techniques are needed to maintain consistency among the local caches (such as to ensure that the local caches do not store different copies of the same shared data). Recent updates to the CXL standard provide cache coherency mechanisms that can be implemented in hardware, by allowing a host device to “snoop” the cache lines of other host devices coupled to the same shared memory, for example, to detect the most recent contents of each cache line. As such, hardware-based cache coherency often requires a significant amount of overhead and generates substantial traffic at larger scales. Aspects of the present disclosure recognize that implementing cache coherency in software can significantly reduce the cost and complexity of a CXL system with shared memory.

(21) Various aspects relate generally to memory management, and more particularly, to techniques for maintaining cache coherency between host devices coupled to a shared memory via a CXL fabric. In some aspects, a memory management system may include a shared memory (SM) manager and multiple SM libraries each associated with a respective host device that is coupled to a shared memory via a CXL fabric. Each SM library communicates with the SM manager to acquire and manage access to the shared memory on behalf of its associated host device. The SM manager creates, configures, and deletes objects in the shared memory responsive to requests from the SM libraries. A memory “object” is a virtual memory allocation that can be managed in software. For example, an object may define an address space in shared memory that can be used for storage by one or more host devices.

(22) In some aspects, an SM library may negotiate read or write “locks” with the SM manager for an object in shared memory, where a read-lock grants the associated host device read access to the object and a write-lock grants the associated host device write access to the object. In some implementations, the SM manager may grant any number of read-locks, but only one write-lock, to any object in shared memory. In other words, an object in shared memory becomes immutable once the host device that created the object relinquishes its write-lock for the object. Immutable objects cannot be modified or rewritten. Any storage in shared memory that is allocated for an immutable object can be written to again only after the object is deleted. In some implementations, the SM

manager may require each host device that has accessed an object in shared memory to flush or otherwise clear its cache before deleting the object.

(23) Particular implementations of the subject matter described in this disclosure can be implemented to realize one or more of the following potential advantages. By creating immutable objects in shared memory, aspects of the present disclosure can ensure cache coherency among host devices coupled to the shared memory via a CXL fabric. Because only the creator of an object can write to the object, and because the data stored in the object cannot be modified after the object is made immutable, the host devices cannot have inconsistent copies of shared data in their local caches. By requiring all host devices that accessed an object in shared memory to flush their caches before the object can be deleted, aspects of the present disclosure can further ensure that none of the host devices have stale or invalid data in their local caches.

(24) FIG. 1 shows a hardware diagram of an example computer system **100** with pooled memory resources. The computer system includes a pool manager **110**, a memory pool **120**, a memory switch **130**, and one or more processing resources **140**. In some aspects, the processing resources **140** may include one or more computing servers.

(25) The memory pool **120** may include any number of memory resources, such as volatile memory, non-volatile memory, or any combination thereof. In some implementations, the memory resources in the memory pool **120** may be distributed across different locations or devices. In the example of FIG. 1, a first portion of the memory pool **120** is disposed on a first server rack **101**, a second portion of the memory pool **120** is disposed on a second server rack **102**, and a third portion of the memory pool **120** is disposed on a third server rack **103**. In some implementations, the number or types of memory resources disposed on each of the server racks **101-103** may vary. For example, the first portion of the memory pool **120** (disposed on the first server rack **101**) may include a different number of memory resources than the second portion of the memory pool **120** (disposed on the second server rack **102**).

(26) The memory switch **130** serves as a memory access interface between the processing resources **140** and the memory pool **120**. More specifically, the memory switch **130** may allow any of the processing resources **140** to write data to, or read data from, any of the memory resources in the memory pool **120** (regardless of their physical locations). For example, the memory switch **130** may facilitate memory access operations (such as read or write access) between processing resources **140** disposed on the first server rack **101** and memory resources disposed on any of the second or third server racks **102** or **103**. In some implementations, the memory switch **130** may form at least part of a Compute Express Link (CXL) fabric conforming to the CXL open standard. For example, the memory switch **130** may be a CXL switch.

(27) The pool manager **110** manages or maintains the memory pool **120**. In some aspects, the pool manager **110** may dynamically provision memory resources in the memory pool **120** to be used by the processing resources **140**. For example, the pool manager **110** may allocate a subset of memory resources from the memory pool **120** to a particular processing resource **140**, when needed, and may deallocate the subset of memory resources when they are no longer needed by the particular processing resource **140**. Aspects of the present disclosure recognize that some pages of memory may become locked (or “pinned”) by a processing resource **140** for long-term or persistent use, which may prevent the processing resource **140** from releasing the entire extent of memory containing the pinned page.

(28) In some implementations, the memory resources in the memory pool **120** may be dynamically allocated on a per-page basis. For example, the pool manager **110** may map the entire memory pool **120** (or a portion thereof) to a group of processing resources **140** and may dynamically allocate (and deallocate) individual pages of memory associated with the memory pool **120** for individual processing resources **140** within the group. As a result, a processing resource **140** may be provisioned a subset of memory pages associated with the memory pool **120**, when needed, and may release (back into the memory pool **120**) individual pages of memory that are no longer in use.

More specifically, the processing resource **140** may continue to maintain any pages of memory that are pinned or otherwise needed for future use.

(29) In some other aspects, the pool manager **110** may allow sharing of memory resources in the memory pool **120**. For example, the pool manager **110** may provision a set of memory resources from the memory pool **120** to be concurrently accessible to multiple processing resources **140** (such as for distributed computing or data processing applications). Still further, in some aspects, the pool manager **110** may transparently “heal” memory that is worn or is otherwise failing. For example, upon detecting one or more failing memory resources (such as indicated by an error rate associated with memory access operations), the pool manager **110** may provision new memory resources in the memory pool **120** to replace the failing memory resources and may proceed to migrate any data stored on the failing memory resources to the newly allocated memory resources.

(30) FIG. 2 shows a block diagram of an example computer system **200** with pooled memory services, according to some implementations. In some implementations, the computer system **200** may be one example of the computer system **100** of FIG. 1. For example, the computer system **200** is shown to include the memory pool **120**, the switch **130**, and the processing resources **140** of FIG. 1. As shown in FIG. 2, the processing resources **140** are coupled to the memory switch **130** via a high-bandwidth, low-latency interconnect **201**. In some implementations, the interconnect **201** may be a Peripheral Component Interconnect Express (PCIe) link.

(31) The computer system **200** further includes a pool manager **210**, a memory machine **220**, and a memory viewer **230**. In some implementations, the pool manager **210** may be one example of the pool manager **110** of FIG. 1. As described with reference to FIG. 1, the pool manager **210** is configured to manage or maintain the memory pool **120**. In some aspects, the pool manager **210** may provide a number of services in connection with managing the memory pool **120** (also referred to as “memory services”). Memory services can be categorized or classified according to various functionality, hardware, software, or performance requirements. Example memory service classes may include memory pooling services, memory sharing services, capacity optimization services, data protection services, security services, and global insights, among other examples. Thus, different classes of memory service may require different configurations of the CXL fabric. For example, memory sharing services may require multiple host devices to be able to communicate with the same memory (or CXL) device. On the other hand, capacity optimization services may require the same host device to be able to communicate with multiple memory (or CXL) devices.

(32) In some implementations, the memory pooling services may include a service for partitioning the memory pool **120**. For example, the pool manager **210** may maintain an allocation table that indicates which memory resources in the memory pool **120** have been allocated to which processing resources **140**. In some other implementations, the memory pooling services may include a service for allocating an unshared memory partition to a system with orchestration of hot add (such as when new physical memory resources are added to the memory pool **120**. In some other implementations, the memory pooling services may include a service for orchestrating hot remove and freeing memory partitions back to the memory pool **120**. For example, the pool manager **210** may deallocate memory resources that are no longer in use by any of the processing resources **140**. Still further, in some implementations, the memory pooling services may include a service for allocating a memory partition to multiple systems in a cluster and orchestrating page-level sub-allocation within that cluster-wide mapped pool (such as described with reference to FIG. 1).

(33) In some implementations, the data protection services may include a service for capturing or acquiring a snapshot of memory. For example, the pool manager **210** may store a copy of the data in one or more memory resources in the memory pool **120** at any given time. In some other implementations, the data protection services may include a service for adding a check signature for each page of memory in a given snapshot. For example, the pool manager **210** may store a checksum that can be used to validate or otherwise detect errors in the snapshot for one or more

pages of memory. In some other implementations, the data protection services may include a service for replicating N-of-M redundancy across independent memory failure domains. For example, the pool manager **210** may store a number (M) of units of data that includes a number (N) of units of original information plus redundancy, where any N data units of the M stored data units can be used to recover the original information in its entirety). Still further, in some implementations, the data protection services may include a service for serializing data to storage. For example, the pool manager **210** may copy data in the memory pool **120** to a data store such as, for example, an HDD or an SSD.

(34) In some implementations, the memory sharing services may include a service for allocating a shared memory partition to multiple systems with orchestration of hot add and configuration of coherency ownership. For example, the pool manager **210** may maintain cache coherency for processing resources **140** that share concurrent access to the same subset of memory resources (such as described with reference to FIG. 1). In some other implementations, the memory sharing services may include a service for naming memory spaces. For example, the pool manager **210** may assign names to shared memory spaces within the memory pool **120** so that the shared memory spaces can be identified by name and mapped to a process address space, thus allowing the data stored in the named memory spaces to be shared by multiple data processes. Still further, in some implementations, the memory sharing services may include a service for cloning snapshots. For example, the pool manager **210** may provide a copy of a snapshot of the memory pool **120** to one or more processing resources **140** to allow the processing resources **140** to quickly bootup to a desired memory state or perform different processes or modifications on the stored data without losing the original “baseline” memory state.

(35) In some implementations, the capacity optimization services may include a service for tiering to compressed memory. For example, the pool manager **210** may move data from the memory pool **120** to compressed memory, such as to optimize for cost, load, or various other factors. In some other implementations, the capacity optimization service may include a service for tiering to storage. For example, the pool manager **210** may move data from the memory pool **120** to a data store, such as to optimize for cost, load, or various other factors. In some implementations, the security services may include a service for access control. For example, the pool manager **210** may prevent processing resources **140** from accessing data in the memory pool **120** to which they are not authorized. In some other implementations, the security services may include a service for managing transport encryption. For example, the pool manager **210** may add point-to-point encryption of data as it is communicated to or from the memory pool **120** (such as from one end of the interconnect **201** to another).

(36) The memory machine **220** may reside with the processing resources **140** to implement various memory services supported by the pool manager **210** and provide hardware application programming interface (API) integration for the processing resources **140**. In some implementations, the memory machine **220** may expand the memory capacity available to the processing resources **140** by providing a software-defined memory pool with intelligent auto-tiering. For example, the memory machine **220** may implement memory snapshot, memory tiering, or memory sharing services for individual processing resources **140**. In some other implementations, the memory machine **220** may accelerate the time-to-discovery for the processing resources **140** by providing transparent checkpointing or other transparent memory services. For example, the memory machine **220** may allow a processing resource **140** to roll back, restore, or clone data from any of the memory resources in the memory pool **120** at any given time. Still further, in some implementations, the memory machine **220** may reduce the cost of cloud services by enabling long-running applications to use low-cost cloud automation instances or by integrating with cloud automation and scheduler to automatically recover from preemptions.

(37) The memory viewer **230** is configured to provide system administrators (or users) with comprehensive insight into the physical topology of the memory pool **120** as well as how

application memory is utilized. In some implementations, the memory viewer **230** may allow system administrators to ensure memory configuration accuracy by providing a real-time view of the memory configuration as well as the bandwidth of each type of memory hardware. For example, such information may be presented in a graphical user interface (GUI) that can be displayed to a user or system administrator. In some other implementations, the memory viewer **230** may help system administrators to optimize their memory configurations for a given workload by providing a heatmap that shows the memory usage patterns of different types of memory resources (with different access latencies). For example, the system administrator may implement memory tiering based on the information provided in the heatmap (such as by moving “cold” data that is infrequently used to lower-cost memory with higher access latencies).

(38) FIG. **3** shows a block diagram of an example memory management system **300**, according to some implementations. In some implementations, the system **300** may be one example of any of the computer systems **100** or **200** of FIGS. **1** and **2**, respectively. For example, the memory management system **300** may support pooling, sharing, or dynamically provisioning external memory for distributed processing resources. In the example of FIG. **3**, the external memory is disposed on a memory appliance **305** and the processing resources are disposed on a number of host computing devices **301-304** communicatively coupled to the memory appliance **305** via a data plane **306**.

(39) The memory appliance **305** may be any device or hardware that aggregates or otherwise provides memory resources for access by an external host. In some aspects, the memory appliance may be a CXL device. With reference to FIGS. **1** and **2**, the memory appliance **305** may be one example of the memory pool **120**. In some implementations, the memory appliance **305** may not have any local processing resources (such as a Type 3 CXL device or memory expander). In some other implementations, the memory appliance **305** also may include local processing resources (such as Type 2 CXL device or a server). In such implementations, the external memory resources provided by the memory appliance **305** may include memory that is locally attached to one or more processing resources residing on the memory appliance **305**.

(40) The host computing devices **301-304** may include any device or hardware having local processing resources (such as a server) that can write data to, or read data from, the external memory provided by the memory appliance **305**. In some aspects, each of the host computing devices **301-304** may represent a root complex of a respective CXL subsystem. With reference to FIGS. **1** and **2**, the host computing devices **301-304** may be examples of the processing resources **140**. The data plane **306** may include various physical or logical interconnects that support the transfer of data between the memory appliance **305** and the host computing devices **301-304**. In some aspects, the data plane **306** may include a CXL fabric that conforms to the CXL open standard. With reference for example to FIGS. **1** and **2**, the data plane **306** may include the memory switch **130**, the interconnect **201**, or any combination thereof.

(41) In some aspects, the memory management system **300** may dynamically configure (and reconfigure) the data plane **306** to support various memory transactions and services between the memory appliance **305** and the host computing devices **301-304**. As described with reference to FIGS. **1** and **2**, example memory services may include, among other examples, memory pooling services (including capacity monitoring and resource provisioning), capacity optimization services (including performance monitoring and tiering), data sharing services (including software-enforced memory coherency), and data protection services (including user software snapshot). Although one memory appliance **305** and four host computing devices **301-304** are depicted in the example of FIG. **3**, in actual implementations the memory management system **300** may dynamically configure the data plane **306** to support memory transactions and services between any number of memory appliances and any number of host computing devices (including multiple CXL systems or multiple CXL subsystems).

(42) The memory management system **300** includes an orchestration component **310**, host agents

321-324, an orchestration user interface (UI) **330**, and endpoint controller **340**. The endpoint controller **340** is configured to manage the memory resources disposed on, or attached to, the memory appliance **305**. In some implementations, the endpoint controller **340** may reside or execute on the memory appliance **305**. For example, the endpoint controller **340** may manage read or write access (by the host computing devices **301-304**) to a memory pool or various logical devices (LDs), including single logical devices (SLDs) and multiple logical devices (MLDs), among other examples.

(43) Each of the host agents **321-324** is configured to provide one or more memory services for a respective one of the host computing devices **301-304**. In some implementations, the host agents **321-324** may reside or execute on the host computing devices **301-304**, respectively, and may communicate with the endpoint controller **340** via the orchestration component **310**. With reference to FIG. 2, each of the host agents **321-324** may be one example or instance of the memory machine **220**. In some implementations, each of the host agents **321-324** may manage multiple components of each supported memory service, including core and data services. Core service components may provide interoperability with other components of the memory management system **300**. Example core service components may include local databases, hardware driver support, or interconnect management software, among other examples. By contrast, data service components may provide non-core functionality. Example data service components may include performance monitoring, external memory provisioning, or software-coherency controls, among other examples.

(44) The orchestration component **310** coordinates interactions between the endpoint controller **340** and the host agents **321-324**. With reference to FIG. 2, the orchestration component **310** may be one example of the pool manager **210**. In some implementations, the orchestration component **310** may manage a global shared state between disparate hosts. For example, the global shared state may include memory provisioning information indicating portions of the external memory resources (if any) that are owned by individual host computing devices **301-304** (such as where each of the host computing devices is associated with a different coherency domain). In some aspects, the orchestration component **310** may reside or execute on one of the host computing devices **301-304**. In some other aspects, the orchestration component **310** may reside or execute on the memory appliance **305**. Still further, in some aspects, the orchestration component **310** may reside or execute on a dedicated host computing device that does not include a host agent or endpoint controller.

(45) The orchestration UI **330** may be any user interface (such as a GUI or command line) that can interact with the orchestration component **310**, the host agents **321-324**, or the endpoint controller **340** to provide external memory orchestration. With reference to FIG. 2, the orchestration UI **330** may be one example of the memory viewer **230**. In some implementations, the orchestration UI **330** may be a web browser-based graphical software suite which communicates with other components of the memory management system **300** via hypertext transfer protocol (HTTP) requests. In some other implementations, the orchestration UI **330** may be a command line terminal software suite which communicates with other components of the memory management system **300** via remote procedure calls (RPCs). In some aspects, the orchestration UI **330** may reside or execute on the same computer system as the orchestration component **310**. In some other aspects, the orchestration UI **330** may reside or execute on a personal computer that is communicatively coupled to the computer system executing the orchestration component **310** (such as via one or more networks).

(46) In some aspects, the memory management system **300** may provide a hardware abstraction layer (HAL) or interface between the memory services and a CXL fabric that couples or otherwise interconnects the host computing devices **301-304** and the memory appliance **305**. For example, the CXL standard defines a “fabric manager” that controls aspects of a CXL system related to binding and management of pooled ports and devices. More specifically, the CXL standard provides a fabric manager API that can be used to configure a CXL fabric. In some implementations, the

orchestration component **310** (or other components of the memory management system **300**) may implement or utilize the fabric manager API to dynamically configure (and reconfigure) the data plane **306** based on the needs or requests of the memory services. For example, the data plane **306** (or CXL fabric) can be used to provision memory resources in the memory appliance **305** for storing data associated with one or more of the host computing devices **301-304**. In other words, the memory management system **300** exposes the features and capabilities of the of the CXL fabric to the memory services supported by the host computing devices **301-304**.

(47) FIG. 4 shows a block diagram of an example host agent **400**, according to some implementations. The host agent **400** is a component of a memory management system that provides one or more memory services (including core services and data services) for a host computing device of a CXL system or subsystem. For example, the host agent **400** may reside or execute on the respective host computing device. In some implementations, the memory management system may be one example of the memory management system **300** of FIG. 3. With reference for example to FIG. 3, the host agent **400** may be one example of any of the host agents **321-324**.

(48) In some aspects, the host agent **400** may provide the memory services based, at least in part, on one or more external memory resources. As used herein, the term “external memory resource” may refer to any memory that is attached to a CXL device (such as the memory appliance **305** of FIG. 3) or is otherwise coupled to the host computing device via a CXL fabric (such as the data plane **306**). In some implementations, the host agent **400** may dynamically configure (and reconfigure) the CXL fabric to facilitate data transactions or other communications between the host computing device and the external memory resources. The host agent **400** includes a host agent controller **410**, a core services component **420**, a system monitoring component **430**, and a data services component **440**. The host agent controller **410** manages interactions between the components **420**, **430**, and **440** of the host agent **400** and other components of the memory management system (such as the orchestration component **310** of FIG. 3) to implement the memory services supported by the host agent **400**.

(49) The core services component **420** enables the host agent **400** to access or communicate with other components of a CXL system (or subsystem). In some implementations, the core services component **420** may include a database management service **422**, an interconnect management service **424**, and a hardware management service **426**. The database management service **422** is configured to manage a local database (on the host computing device) for storing data and other information associated with one or more memory services. The interconnect management service **424** is configured to create and manage pathways for routing data and other information through the CXL fabric. The hardware management service **426** is configured to manage various hardware drivers on the host computing device so that the host agent **400** can operate other hardware components of the CXL system. The core services component **420** may configure the database, interconnects, or hardware drivers responsive to control signals **402** from the host agent controller **410**.

(50) The system monitoring component **430** is configured to monitor various aspects of the host computing device or the CXL system (or subsystem). In some implementations, the system monitoring component **430** may include a performance monitoring service **432**, a capacity monitoring service **434**, and a health monitoring service **436**. The performance monitoring service **432** is configured to monitor various performance metrics associated with the host computing device or the CXL system (such as a bandwidth, speed, latency, or power consumption of data transactions). The capacity monitoring service **434** is configured to monitor a storage capacity of any memory (or other data storage) available to the host computing device (including locally-attached memory or external memory resources). The health monitoring service **436** is configured to monitor a health of various hardware components of the host computing device (such as memory failures or other device or component failures). The system monitoring component **430** may report

the monitored conditions to the host agent controller **410** via system updates **404**.

(51) The data services component **440** is configured to provide non-core services related to the management of data. In some implementations, the data services component **440** may include a memory snapshot service **442**, a memory sharing service **444**, a memory tiering service **446**, and a memory elasticity service **448**. The memory snapshot service **442** is configured to capture or acquire a snapshot of data stored in memory (such as locally-attached memory or external memory). The memory sharing service **444** is configured to manage or maintain one or more cache coherency states associated with any data stored in memory shared with other host computing devices. The memory tiering service **446** is configured to manage the movement or storage of data to different storage media based on various memory tiers. The memory elasticity service **448** is configured to dynamically allocate memory resources from a memory pool to the host computing device (and release the memory resources back into the memory pool). The data services component **440** may perform various actions associated with the data services responsive to service requests **406** from the host agent controller **410**.

(52) In the example of FIG. 4, the data services component **440** is shown to provide four data services for the host computing device. In some other implementations, the data services component **440** may provide fewer or more data services than those depicted in FIG. 4. As described with reference to FIG. 2, memory snapshot services, memory sharing services, memory tiering services, and memory elasticity services belong to different classes of memory services. More specifically, each of the data services may require the CXL fabric to be configured in a particular way (such as to provision specific memory resources suitable for the data service). In some aspects, the host agent **400** may configure the CXL fabric based, at least in part, on the class of memory service associated with each data service provide by the data services component **440**.

(53) In some implementations, to support memory snapshot services, the host agent controller **410** may produce one or more control signals **402** instructing the interconnect management service **424** to establish a data path, via the CXL fabric, from the host computing device to an external memory suitable for storing a snapshot of any data stored in locally-attached memory. The host agent controller **410** may further produce one or more service requests **406** instructing the memory snapshot service **442** to provision at least a portion of the external memory to be used for storing the snapshots of the locally-attached memory. Thereafter, the memory snapshot service **442** may capture or acquire snapshots of the data stored in the locally-attached memory (such as in response to various trigger conditions) and store the snapshots in the external memory.

(54) In some implementations, to support memory sharing services, the host agent controller **410** may produce one or more control signals **402** instructing the interconnect management service **424** to establish a data path, via the CXL fabric, from the host computing device to an external memory suitable for storing data that can be shared among multiple host computing devices. The host agent controller **410** may further produce one or more service requests **406** instructing the memory sharing service **444** to provision at least a portion of the external memory to be used for sharing data between the host computing devices. Thereafter, the memory sharing service **444** may manage one or more cache coherency states associated with the data stored in the shared memory (such as in response to memory access requests initiated by the local host computing device or another host computing device having access to the shared memory).

(55) In some implementations, to support memory tiering services, the host agent controller **410** may produce one or more control signals **402** instructing the interconnect management service **424** to establish a data path, via the CXL fabric, from the host computing device to an external memory associated with a different memory tier than the locally-attached memory. The host agent controller **410** may further produce one or more service requests **406** instructing the memory tiering service **446** to provision at least a portion of the external memory to be used for offloading data assigned to a lower (or higher) memory tier than that of the locally-attached memory. Thereafter, the memory tiering service **446** may move data between the locally-attached memory and the external memory

based on various triggers that assign (or reassign) the data to different memory tiers.

(56) In some implementations, to support elastic memory services, the host agent controller **410** may produce one or more control signals **402** instructing the interconnect management service **424** to establish a data path, via the CXL fabric, from the host computing device to an external memory associated with a memory pool. The host agent controller **410** may further produce one or more service requests **406** instructing the memory elasticity service **448** to provision at least a portion of the external memory to be used by one or more virtual machines (or tasks) executing on the host computing device. Thereafter, the memory elasticity service **448** may dynamically allocate (or release) resources from the memory pool based on the memory requirements of each system workload.

(57) Aspects of the present disclosure recognize that updates associated with the data services may require the CXL fabric to be reconfigured (such as to provision new memory resources, deprovision existing memory resources, or otherwise rebalance or optimize memory usage). In some aspects, the host agent **400** may dynamically reconfigure the CXL fabric based on updates received from the host computing device or other device coupled to the CXL fabric (such as a CXL device or another host computing device). Example suitable updates include, among other examples, manual inputs received from a user of the host computing device (such as via the orchestration UI **330** of FIG. 3), performance-related events indicating a change in the performance of the CXL system, capacity-related events indicating a change in the storage capacity associated with the CXL system, health-related events indicating a change in the health of the CXL system, and various requests initiated by applications executing on the host computing device (such as the launch of a hypervisor or scheduler integration).

(58) In some aspects, the host agent **400** may reconfigure the CXL fabric based on updates associated with memory snapshot services. For example, the host agent controller **410** may produce one or more control signals **402** instructing the interconnect management service **424** to tear down one or more existing data paths through the CXL fabric or to establish a new data path from the host computing device to another external memory or storage (such as an HDD or an SSD). Alternatively, or in addition, the host agent controller **410** may produce one or more service requests **406** instructing the memory snapshot service **442** to deprovision one or more external memory resources currently allocated to the host computing device or to provision other external memory resources to be used for storing snapshots of the locally-attached memory. In some implementations, the host agent controller **410** may produce one or more service requests **406** instructing the memory snapshot service **442** to offload a snapshot from external memory to a more permanent storage solution (such as HDD or SSD).

(59) In some aspects, the host agent **400** may reconfigure the CXL fabric based on updates associated with memory sharing services. For example, the host agent controller **410** may produce one or more control signals **402** instructing the interconnect management service **424** to tear down one or more existing data paths through the CXL fabric or to establish a new data path from the host computing device to another external memory. Alternatively, or in addition, the host agent controller **410** may produce one or more service requests **406** instructing the memory sharing service **444** to deprovision one or more external memory resources currently allocated to the host computing device or to provision other external memory resources to be used for storing data shared by multiple host computing devices. In some implementations, the host agent controller **410** may produce one or more service requests **406** instructing the memory sharing service **444** to change the cache coherency states of one or more cache lines associated with any of the shared data.

(60) In some aspects, the host agent **400** may reconfigure the CXL fabric based on updates associated with memory tiering services. For example, the host agent controller **410** may produce one or more control signals **402** instructing the interconnect management service **424** to tear down one or more existing data paths through the CXL fabric or to establish a new data path from the

host computing device to another external memory or storage (such as an HDD or an SSD). Alternatively, or in addition, the host agent controller **410** may produce one or more service requests **406** instructing the memory tiering service **446** to deprovision one or more external memory resources currently allocated to the host computing device or to provision other external memory resources to be used for offloading data assigned to various memory tiers. In some implementations, the newly-provisioned memory resources may be associated with an even lower (or higher) memory tier than that of any memory resources currently allocated to the host computing device.

(61) In some aspects, the host agent **400** may reconfigure the CXL fabric based on updates associated with elastic memory services. For example, the host agent controller **410** may produce one or more control signals **402** instructing the interconnect management service **424** to tear down one or more existing data paths through the CXL fabric or to establish a new data path from the host computing device to another external memory associated with the memory pool. Alternatively, or in addition, the host agent controller **410** may produce one or more service requests **406** instructing the memory elasticity service **448** to deprovision one or more external memory resources currently allocated to the host computing device or to provision other external memory resources from the memory pool. In some implementations, the elastic memory service **448** may dynamically expand the amount of memory available to the host computing device, from the memory pool, in response to increased memory usage or demand by virtual machines and may release memory back into the memory pool as usage or demand subsides.

(62) FIG. 5 shows another block diagram of an example memory management system **500**, according to some implementations. The memory management system **500** includes a shared memory (SM) manager **510** and multiple SM libraries **521-523** that reside (or execute) on host computing devices **501-503**, respectively. In some implementations, the memory management system **500** may be one example of the memory management system **300** of FIG. 3. With reference to FIG. 3, the SM manager **510** may be one example of the orchestration component **310** and each of the SM libraries **521-523** may be one example of any of the host agents **321-324**. In the example of FIG. 5, the SM manager **510** is shown to reside (or execute) on the host device **503**. However, in actual implementations, the SM manager **510** may reside on any of the host devices **501-503** or another device that does not include an SM library.

(63) Each of the host devices **501-503** includes a processor (labeled “CPU”) and a local cache (labeled “CPU cache”). In some implementations, the host devices **501-503** may be coupled to a shared memory **530** via a CXL fabric **540**. With reference to FIG. 3, each of the host devices **501-503** may be one example of any of the host devices **301-304**, the shared memory **530** may be one example of the memory appliance **305**, and the CXL fabric **540** may be one example of the data plane **306**. The memory management system **500** provides software-based cache coherency management between the host devices **501-503**. More specifically, the SM manager **510** may coordinate with the SM libraries **521-523** to detect and execute the required access semantics (such as which host device is requesting access to what portion of the shared memory **530** at any given time) in a manner that maintains coherency among the local caches.

(64) In some aspects, the memory management system **500** may implement a page-based share semantic detection mechanism to detect and control access to pages of the shared memory **530** (or other granularities of memory). In such aspects, the memory management system **500** may act as a logic gate that intercepts requests from the host devices **501-503** to access one or more pages of the shared memory **530** (in hardware) and selectively grants each memory access request only when the requested access would not create cache coherency issues with any of the other host devices. For example, the memory management system **500** may remove read and write access to such pages of the shared memory **530** and may clear the caches of one or more the host devices **501-503** before granting the requested memory access to a particular host device.

(65) In some implementations, the SM manager **510** may maintain a cache coherency table that

(71) In some other aspects, the memory management system **500** may implement a procedure-based share semantic detection mechanism to detect and control access to objects in the shared memory **530**. In such aspects, the memory management system **500** may create, manage, and delete objects in the shared memory **530** responsive to requests received from the host devices **501-503**. In some implementations, the memory management system **500** may detect the requested access semantics by intercepting existing sharing-semantic procedure calls (such as RPCs). In some other implementations, the memory management system **500** may detect the requested access semantics by substituting or replacing existing sharing-semantic procedure calls with new shared memory application interfaces (APIs).

(72) In some implementations, the SM manager **510** may provide shared object management services for the host devices **501-503**. Example suitable shared object management services may include object creation, configuration, and deletion, among other examples. In some implementations, the SM manager **510** may further provide configuration management services for the host devices **501-503**. For example, the SM manager **510** may manage multiple shared memories (not shown for simplicity) each having its own configuration. In such implementations, the configuration management service may manage a respective configuration for each of the shared memories.

(73) The SM libraries **521-523** provide shared memory services for individual host devices **501-503**, respectively. In some implementations, the SM libraries **521-523** may enhance existing RPCs that express sharing semantics to encapsulate low-level operations for managing coherency or supporting other services provided by the SM manager **510**. In some other implementations, the SM libraries **521-523** may provide new low-level APIs associated with shared memory services. Example shared memory (SM) API requests may include: an “SM_connect” request to connect with the SM manager **510**; an “SM_disconnect” request to disconnect from the SM manager **510**; an “SM_create” request to create an object in the shared memory **530** (acquiring write-access to the object on behalf of the host device); an “SM_seal” request to seal an object in the shared memory **530** (relinquishing write-access to the object on behalf of the host device); an “SM_get” request to get an object in the shared memory **530** (acquiring read-access to the object on behalf of the host device); an “SM_release” request to release an object in the shared memory **530** (relinquishing read-access to the object on behalf of the host device); an “SM_delete” request to delete an object in the shared memory **530**; and an “SM_subscribe” request to subscribe to notifications indicating when objects in the shared memory **530** are sealed.

(74) In some implementations, the SM manager **510** may manage lock services for each of the host devices **501-503**. For example, when an application (“App”) executing on a particular host device requires access to an object in the shared memory **530**, the application must first acquire a “lock” (such as a read-lock or a write-lock) for the shared object. The type of lock granted to a particular host device governs the access rights of that host device. For example, a write-lock grants a host device permission to write to an object in the shared memory **530** until the host device relinquishes the write-lock, whereas a read-lock grants a host device permission to read from an object in the shared memory **530** until the host device relinquishes the read-lock. Host devices may be granted exclusive write-locks, shared read-locks, or any combination thereof. The SM manager **510** manages the locks for the host devices **501-503** and maintains a mapping of the locks granted to each host device for each object in the shared memory **530**.

(75) In some implementations, each of the SM libraries **521-523** may provide individual lock services for its respective host device. For example, when an application executing on a particular host device requires access to the shared memory **530**, the SM library residing on that host device may negotiate with the SM manager **510** to acquire the necessary lock associated with the requested access type (such as read or write access) and map the shared memory **530** with the granted access type. In some implementations, each of the SM libraries **521-523** may further communicate with the SM manager **510** to maintain cache coherency and access synchronization

(or any other configurations associated with the shared memory **530**) across the various host devices **501-503**.

(76) In some aspects, the SM manager **510** may grant only one write-lock for any object in the shared memory **530**. More specifically, only the host device that created an object is granted a write-lock for the object (upon its creation). In some implementations, the SM libraries **521-523** may perform write-through operations (using CPU non-temporal instructions) on behalf of the host devices **501-503**, respectively, to write data directly to the shared memory **530** (rather than writing to a write-back cache followed by a cache flush). When the host device has finished writing to an object, the host device relinquishes its write-lock for that object. In some implementations, an object may become immutable once a host device relinquishes its write-lock for the object. Immutable objects cannot be modified or written to. Thus, any pages of the shared memory **530** that are allocated for an immutable object can be rewritten only after the object is deleted.

(77) To avoid cache coherency issues, the SM manager **510** may wait for a write-lock granted for a particular object to be released or relinquished before granting any read-lock requests for that object. Waiting requests may be queued and processed by the SM manager **510** in order of arrival. In some aspects, the SM manager **510** may grant multiple read-locks, concurrently, for the same object in the shared memory **530**. When a host device has finished reading from an object, the host device relinquishes its read-lock for that object. In some aspects, the SM manager **510** may block or deny any further read-lock requests for an object after receiving a request to delete the object. Further, the SM manager **510** may wait for all existing read-locks granted for a particular object to be released or relinquished before deleting that object from the shared memory **530**.

(78) In some implementations, the memory management system **500** may flush the local caches of one or more of the host devices **501-503** prior to deleting an object from the shared memory **530**. For example, the SM libraries **521-523** may track whether the host devices **501-503**, respectively, have accessed any objects in the shared memory **530** and may flush the caches for their respective host devices when any previously-accessed objects are about to be deleted. In some implementations, each of the SM libraries **521-523** may coordinate with the SM manager **510** to ensure that all locks associated with a shared object have been released before the object is deleted from the shared memory **530**. In some implementations, each of the SM libraries **521-523** also may provide an API for applications executing on its respective host device to flush the local cache on demand (or when necessary).

(79) In some aspects, each of the SM libraries **521-523** may include a memory allocator for allocating shared memory **530** to its respective host device. For example, the memory allocator may provide GNU C Library (glibc) APIs, such as “malloc,” “free,” and “realloc.” In some implementations, the memory allocator may further provide a special memory barrier function which guarantees that the memory written before the function call is visible to other host devices. In some implementations, the memory allocator may be initialized from memory that has already been used. In some implementations, the shared memory **530** may be subdivided into multiple segments, where each segment can be managed and accessed individually (allowing parallel access to the shared memory **530**). In such implementations, each of the SM libraries **521-523** may include a respective instance of the memory allocator for each segment of the shared memory **530**.

(80) In some implementations, each of the SM libraries **521-523** may provide an RPC service over the shared memory **530**. For example, the RPC service may include RPC APIs that are built on top of the memory allocator and are easier to use. Still further, in some aspects, each of the SM libraries **521-523** may be configured to implement tiering services to improve or optimize memory performance. As described with reference to FIGS. **1** and **4**, example suitable tiering operations may include tiering to compressed memory and tiering to storage, among other examples.

(81) In some implementations, the SM manager **510** and SM libraries **521-523** may improve the speed of shared memory access by mapping the address space of the shared memory **530** to the same virtual addresses across all processes associated with the host devices **501-503**. As a result,

the host devices **501-503** may maintain consistent pointer values to the shared memory **530**. In some other implementations, the SM manager **510** and SM libraries **521-523** may improve the speed of shared memory access by allowing objects in the shared memory **530** to only reference other objects in the shared memory **530**. In other words, objects in the shared memory **530** may not reference any other objects outside the shared memory **530**.

(82) FIGS. **6A-6E** show a series of operations **600-640**, respectively, that can be performed by a memory management system, according to some implementations. The memory management system is configured to provide memory sharing services for a number of host computing devices (host device A, host device B, and host device C) coupled to a shared memory **650** via a CXL fabric (not shown for simplicity). The memory management system includes an SM manager **604** and a number of SM libraries **601-603** associated with the host devices A—C, respectively. In some implementations, the memory management system may be one example of the memory management system **500** of FIG. **5**. With reference to FIG. **5**, the SM manager **604** may be one example of the SM manager **510** and each of the SM libraries **601-603** may be one example of any of the SM libraries **521-523**.

(83) With reference to FIG. **6A**, the SM library **602** receives an SM_create request from host device B and negotiates with the SM manager **604** to create an object **606** in the shared memory **650** based on the SM_create request. For example, the SM manager **604** may allocate, via the CXL fabric, a range of physical memory addresses (or pages) in the shared memory **650** for the object **606** as long as the memory space does not overlap with another object in the shared memory **650**. Upon creating the object **606** in the shared memory **650**, the SM manager **604** grants or otherwise assigns a write-lock **607** for the object **606** to host device B. The SM manager **604** creates a mapping for the object **606** indicating that a write-lock is assigned to host device B (WL=B) and that no read-locks are currently assigned (RL=none). The SM library **602** also maps a write-lock (WL) to the object **606** and creates a reminder indicating that a cache flush is required (CFR) when the object **606** is deleted. The SM libraries **601** and **603** receive SM_subscribe requests from host devices A and C, respectively, and request to be notified by the SM manager **604** when the object **606** is sealed.

(84) With reference to FIG. **6B**, the SM library **602** receives an SM_seal request from host device B and instructs the SM manager **604** to seal the object **606** in the shared memory **650** in response to receive the SM_seal request. For example, the SM manager **604** may seal the object **606** in the shared memory **650** by updating the mapping to indicate that the object **606** is immutable or that no further write-locks may be granted for the object **606** (WL=immutable). The SM manager **604** notifies the SM libraries **601** and **603** when the object **606** has been sealed (responsive to the SM_subscribe requests of FIG. **6A**). The SM libraries **601** and **603** receive SM_get requests from host devices A and C, respectively, and negotiate with the SM manager **604** to acquire read-locks for the object **606**. For example, the SM manager **604** may grant a read-lock request as long as no write-locks are currently mapped to the object **606**. In the example of FIG. **6B**, the SM manager **604** grants concurrent read-locks **611** and **613** to host devices A and C, respectively, and updates the mapping to indicate that read-locks are assigned to host devices A and C (RL=A,C). Each of the SM libraries **601** and **603** also maps a read-lock (RL) to the object **606** and creates a reminder indicating that a cache flush is required (CFR) when the object **606** is deleted.

(85) With reference to FIG. **6C**, the SM library **603** receives an SM_release request from host device C and relinquishes its read-lock for the object **606** in the shared memory **650**. In response, the SM manager **604** updates its mapping to indicate that only host device A is currently assigned a read-lock for the object **606** (RL=A). The SM library **602** receives an SM_delete request from host device B and negotiates with the SM manager **604** to delete the object **606** from the shared memory **650**. As described with reference to FIG. **5**, the SM manager **604** may delete a shared object only when all host devices (that accessed the object **606**) have released their respective locks for the object. Because host device A has not yet relinquished its read-lock, the SM manager **604** does not

delete the object **606** from the shared memory **650**. Instead, the SM manager **604** queues an object delete request **622** and blocks or denies any further read-lock requests for the object **606**.

(86) With reference to FIG. **6D**, the SM library **601** receives an SM_release request from host device A and relinquishes its read-lock for the object **606** in the shared memory **650**. In response, the SM manager **604** updates its mapping to indicate that no host devices are currently assigned a read-lock for the object **606** (RL=none). Because all read-locks (and write-locks) for the object **606** have been released, the SM manager **604** may process the object delete request **622** received from host device B (in FIG. **6C**). For example, the SM manager **604** may provide a notification **632** to each of the SM libraries **601-603** indicating that the object **606** is about to be deleted from the shared memory **650**. Each of the SM libraries **601-603** responds to the notification **632** by checking whether its respective host device has a cache flush request (CFR) mapped to the object **606** (indicating that the host device has previously accessed the object **606**) and flushing the local cache of the host device if a CFR is detected. In the example of FIG. **6D**, each of the host devices A-C has a CFR mapped to the object **606**. Thus, the SM libraries **601-603** may flush the caches of the host devices A-C, respectively, in response to the object delete notification **632**.

(87) With reference to FIG. **6E**, each of the SM libraries **601-603** provides a report **642** to the SM manager **604** indicating that the cache of its associated host device is clear of any shared data that may be stored in the object **606**. For example, each of the SM libraries **601-603** may provide the report **642** to the SM manager **604** responsive to flushing the cache of its associated host device. In some implementations, the SM manager **604** may require a report from every host device coupled to the shared memory **650** (regardless of whether the host device had to perform a cache flush operation) before deleting the object **606** from the shared memory **650**. After receiving the report **642** from all of the SM libraries **601-603**, the SM manager **604** may proceed to delete the object **606** from the shared memory **650**. For example, the SM manager **604** may free or return the physical memory addresses (or pages) associated with the object **606** back to the shared memory **650** so that such memory resources can be used by other processes or tasks.

(88) FIG. **7** shows another block diagram of an example memory management system **700**, according to some implementations. The memory management system **700** is configured to provide shared memory services over a CXL system. The CXL system may include one or more host computing devices and one or more CXL devices coupled via a CXL fabric. In some implementations, the memory management system **700** may be one example of any of the memory management systems **300** or **500** of FIGS. **3** and **5**, respectively. With reference for example to FIG. **5**, each host computing device may be one example of any of the host devices **501-503** and the CXL device may be one example of the shared memory **530**, and the CXL fabric may be one example of the CXL fabric **540**.

(89) The memory management system **700** includes a device interface **710**, a processing system **720**, and a memory **730**. The device interface **710** is configured to communicate with one or more devices associated with the CXL system (such as the host computing devices or the CXL devices). For example, the device interface **710** may include a host device interface (I/F) **712** configured to communicate with a host computing device of a CXL subsystem and a CXL device interface (I/F) **714** configured to communicate with a CXL device of the CXL subsystem. In some implementations, the host device I/F **712** may receive, from a host device, requests to acquire and relinquish a write-lock for an object (such as the SM_create and SM_seal requests of FIG. **5**).

(90) The memory **730** may include a data store **732** configured to store data and other information associated with the shared memory services. The memory **730** also may include a non-transitory computer-readable medium (including one or more nonvolatile memory elements, such as EPROM, EEPROM, Flash memory, or a hard drive, among other examples) that may store at least the following software (SW) modules: an object creation SW module **734** to create an object in the memory shared by the plurality of host computing devices over the CXL fabric; a lock management SW module **736** to map a write-lock to the object in the shared memory responsive to creating the

object, where the write-lock grants a first host computing device of the plurality of host computing devices exclusive access for writing to the object; and an object sealing SW module **738** to seal the object in the shared memory so that the object is immutable based on receiving a request to relinquish the write-lock.

Each software module includes instructions that, when executed by the processing system **720**, causes the memory management system **700** to perform the corresponding functions.

(91) The processing system **720** may include any suitable one or more processors capable of executing scripts or instructions of one or more software programs stored in the memory management system **700** (such as in the memory **730**). For example, the processing system **720** may execute the object creation SW module **734** to create an object in the memory shared by the plurality of host computing devices over the CXL fabric. The processing system **720** also may execute the lock management SW module **736** to map a write-lock to the object in the shared memory responsive to creating the object, where the write-lock grants a first host computing device of the plurality of host computing devices exclusive access for writing to the object. Further, the processing system **720** may execute the object sealing SW module **738** to seal the object in the shared memory so that the object is immutable based on receiving a request to relinquish the write-lock.

(92) FIG. **8** shows an illustrative flowchart depicting an example operation **800** for memory management, according to some implementations. In some implementations, the example operation **800** may be performed by a memory management system such as any of the memory management systems **300**, **500**, or **700** of FIGS. **3**, **5**, and **7**, respectively.

(93) The memory management system creates an object in a memory shared by a plurality of host computing devices via a CXL fabric (**810**). In some implementations, each host computing device of the plurality of host computing devices may represent a root complex of a respective CXL subsystem. The memory management system maps a write-lock to the object in the shared memory responsive to creating the object, where the write-lock grants a first host computing device of the plurality of host computing devices exclusive access for writing to the object (**820**). The memory management system further receives a request from the first host computing device to relinquish the write-lock (**830**). The memory management system seals the object in the shared memory so that the object is immutable based on the request to relinquish the write-lock (**840**).

(94) In some aspects, the memory management system may further receive a request from the first host computing device to write data to the object in the shared memory prior to relinquishing the write-lock and may perform a write-through operation that writes the data directly to the shared memory, via one or more CXL protocols, without modifying a cache associated with the first host computing device.

(95) In some aspects, the memory management system may further receive a request from a second host computing device of the plurality of host computing devices to access the object in the shared memory, determine whether one or more locks, including a write-lock or a read-lock, are mapped to the object, and control access to the object by the second host computing device based at least in part on whether one or more locks are mapped to the object. In some implementations, the controlling of access to the object may include preventing the second host computing device from accessing the object while the write-lock is mapped to the object.

(96) In some aspects, the request to access the object in the shared memory may be associated with a request to read from the object. In some implementations, the controlling of access to the object may include mapping a first read-lock to the object based on the access request when no locks are mapped to the object, where the read-lock grants the second host computing device access for reading from the object in the shared memory. In some other implementations, the controlling of access to the object may include mapping a first read-lock to the object based on the access request when one or more second read-locks are mapped to the object, where the first read-lock grants the second host computing device access for reading from the object in the shared memory.

(97) In some aspects, the memory management system may further receive a request to delete the object in the shared memory, block any requests to access the object received after the request to delete the object, and delete the object from the shared memory when no locks are mapped to the object. In some implementations, the deleting of the object may include determining whether each host computing device of the plurality of host computing devices has previously accessed the object and flushing a respective cache associated with any host computing device of the plurality of host computing devices that has previously accessed the object, where the object is deleted only after all the caches associated with any host computing device of the plurality of host computing devices that has previously accessed the object have been flushed.

(98) FIG. 9 shows another illustrative flowchart depicting an example operation 900 for memory management, according to some implementations. In some implementations, the example operation 900 may be performed by a memory management system such as any of the memory management systems 300, 500, or 700 of FIGS. 3, 5, and 7, respectively.

(99) The memory management system receives a request to delete an object in a memory shared by a plurality of host computing devices (910). The memory management system determines whether each host computing device of the plurality of host computing devices has ever accessed the object (920). The memory management system further flushes a respective cache associated with any host computing device of the plurality of host computing devices that has ever accessed the object (930). The memory management system deletes the object from the shared memory after flushing all the caches associated with any host computing device of the plurality of host computing devices that has ever accessed the object (940).

(100) In some aspects, the memory management system may further block any requests to access the object received after the request to delete the object and determine whether one or more locks, including a write-lock or a read-lock, are mapped to the object, the object being deleted when no locks are mapped to the object. In some implementations, each host computing device of the plurality of host computing devices may represent a root complex of a respective CXL subsystem and the memory may be shared by the plurality of host computing devices via a CXL fabric.

(101) Those of skill in the art will appreciate that information and signals may be represented using any of a variety of different technologies and techniques. For example, data, instructions, commands, information, signals, bits, symbols, and chips that may be referenced throughout the above description may be represented by voltages, currents, electromagnetic waves, magnetic fields or particles, optical fields or particles, or any combination thereof.

(102) Further, those of skill in the art will appreciate that the various illustrative logical blocks, modules, circuits, and algorithm steps described in connection with the aspects disclosed herein may be implemented as electronic hardware, computer software, or combinations of both. To clearly illustrate this interchangeability of hardware and software, various illustrative components, blocks, modules, circuits, and steps have been described above generally in terms of their functionality. Whether such functionality is implemented as hardware or software depends upon the particular application and design constraints imposed on the overall system. Skilled artisans may implement the described functionality in varying ways for each particular application, but such implementation decisions should not be interpreted as causing a departure from the scope of the disclosure.

(103) The methods, sequences or algorithms described in connection with the aspects disclosed herein may be embodied directly in hardware, in a software module executed by a processor, or in a combination of the two. A software module may reside in RAM memory, flash memory, ROM memory, EPROM memory, EEPROM memory, registers, hard disk, a removable disk, a CD-ROM, or any other form of storage medium known in the art. An exemplary storage medium is coupled to the processor such that the processor can read information from, and write information to, the storage medium. In the alternative, the storage medium may be integral to the processor.

(104) In the foregoing specification, embodiments have been described with reference to specific

examples thereof. It will, however, be evident that various modifications and changes may be made thereto without departing from the broader scope of the disclosure as set forth in the appended claims. The specification and drawings are, accordingly, to be regarded in an illustrative sense rather than a restrictive sense.

Claims

1. A method of memory management, comprising: creating a virtual memory object in a memory shared by a plurality of host computing devices via a Compute Express Link (CXL) fabric; mapping a write-lock to the virtual memory object in the shared memory responsive to creating the virtual memory object, the write-lock granting a first host computing device of the plurality of host computing devices exclusive access for writing to the virtual memory object; receiving a request from the first host computing device to relinquish the write-lock; and sealing, in response to receiving the request to relinquish the write-lock, the virtual memory object in the shared memory so that the virtual memory object is immutable.
2. The method of claim 1, wherein each host computing device of the plurality of host computing devices represents a root complex of a respective CXL subsystem.
3. The method of claim 1, further comprising: receiving a request from the first host computing device to write data to the virtual memory object in the shared memory prior to relinquishing the write-lock; and performing a write-through operation that writes the data directly to the shared memory, via one or more CXL protocols, without modifying a cache associated with the first host computing device.
4. The method of claim 1, further comprising: receiving a request from a second host computing device of the plurality of host computing devices to access the virtual memory object in the shared memory; determining whether one or more locks, including a write-lock or a read-lock, are mapped to the virtual memory object; and controlling access to the virtual memory object by the second host computing device based at least in part on whether one or more locks are mapped to the virtual memory object.
5. The method of claim 4, wherein the controlling of access to the virtual memory object comprises: preventing the second host computing device from accessing the virtual memory object while the write-lock is mapped to the virtual memory object.
6. The method of claim 4, wherein the request to access the virtual memory object in the shared memory is associated with a request to read from the virtual memory object.
7. The method of claim 6, wherein the controlling of access to the virtual memory object comprises: mapping a first read-lock to the virtual memory object based on the access request when no locks are mapped to the virtual memory object, the read-lock granting the second host computing device access for reading from the virtual memory object in the shared memory.
8. The method of claim 6, wherein the controlling of access to the virtual memory object comprises: mapping a first read-lock to the virtual memory object based on the access request when one or more second read-locks are mapped to the virtual memory object, the first read-lock granting the second host computing device access for reading from the virtual memory object in the shared memory.
9. The method of claim 1, further comprising: receiving a request to delete the virtual memory object in the shared memory; blocking any requests to access the virtual memory object received after the request to delete the virtual memory object; and deleting the virtual memory object from the shared memory when no locks are mapped to the virtual memory object.
10. The method of claim 9, wherein the deleting of the virtual memory object comprises: determining whether each host computing device of the plurality of host computing devices has previously accessed the virtual memory object; and flushing a respective cache associated with any host computing device of the plurality of host computing devices that has previously accessed the

virtual memory object, the virtual memory object being deleted only after all the caches associated with any host computing device of the plurality of host computing devices that has previously accessed the virtual memory object have been flushed.

11. A memory management system comprising: a processing system; and a memory storing instructions that, when executed by the processing system, causes the memory management system to: create a virtual memory object in a memory shared by a plurality of host computing devices via a Compute Express Link (CXL) fabric; map a write-lock to the virtual memory object in the shared memory responsive to creating the virtual memory object, the write-lock granting a first host computing device of the plurality of host computing devices exclusive access for writing to the virtual memory object; receive a request from the first host computing device to relinquish the write-lock; and seal, in response to receiving the request to relinquish the write-lock, the virtual memory object in the shared memory so that the virtual memory object is immutable.

12. The memory management system of claim 11, wherein each host computing device of the plurality of host computing devices represents a root complex of a respective CXL subsystem.

13. The memory management system of claim 11, wherein execution of the instructions further causes the memory management system to: receive a request from the first host computing device to write data to the virtual memory object in the shared memory prior to relinquishing the write-lock; and perform a write-through operation that writes the data directly to the shared memory, via one or more CXL protocols, without modifying a cache associated with the first host computing device.

14. The memory management system of claim 11, wherein execution of the instructions further causes the memory management system to: receive a request from a second host computing device of the plurality of host computing devices to access the virtual memory object in the shared memory; determine whether one or more locks, including a write-lock or a read-lock, are mapped to the virtual memory object; and control access to the virtual memory object by the second host computing device based at least in part on whether one or more locks are mapped to the virtual memory object.

15. The memory management system of claim 14, wherein the controlling of access to the virtual memory object comprises: preventing the second host computing device from accessing the virtual memory object while the write-lock is mapped to the virtual memory object.

16. The memory management system of claim 14, wherein the request to access the virtual memory object in the shared memory is associated with a request to read from the virtual memory object, the controlling of access to the virtual memory object comprising: mapping a first read-lock to the virtual memory object based on the access request when one or more second read-locks are mapped to the virtual memory object or when no locks are mapped to the virtual memory object, the first read-lock granting the second host computing device access for reading from the virtual memory object in the shared memory.

17. The memory management system of claim 11, wherein execution of the instructions further causes the memory management system to: receive a request to delete the virtual memory object in the shared memory; block any requests to access the object received after the request to delete the virtual memory object; determine whether each host computing device of the plurality of host computing devices has previously accessed the virtual memory object; flush a respective cache associated with any host computing device of the plurality of host computing devices that has previously accessed the virtual memory object; and delete the virtual memory object from the shared memory when no locks are mapped to the virtual memory object and after all the caches associated with any host computing device of the plurality of host computing devices that has previously accessed the virtual memory object have been flushed.

18. A method of memory management, comprising: receiving a request to delete an object in a memory shared by a plurality of host computing devices; determining whether each host computing device of the plurality of host computing devices has ever accessed the object; flushing a respective

cache associated with any host computing device of the plurality of host computing devices that has ever accessed the object; and deleting the object from the shared memory after flushing all the caches associated with any host computing device of the plurality of host computing devices that has ever accessed the object.

19. The method of claim 18, further comprising: blocking any requests to access the object received after the request to delete the object; and determining whether one or more locks, including a write-lock or a read-lock, are mapped to the object, the object being deleted when no locks are mapped to the object.

20. The method of claim 18, wherein each host computing device of the plurality of host computing devices represents a root complex of a respective Compute Express Link (CXL) subsystem and the memory is shared by the plurality of host computing devices via a CXL fabric.
