

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12399845
Kind Code	B2
Date of Patent	August 26, 2025
Inventor(s)	Awasthi; Aditya et al.

Memory organization and access for efficient matrix operations

Abstract

Certain aspects of the present disclosure provide techniques and apparatus for efficiently accessing memory in a computing system. An example method includes organizing a plurality of physical memory banks having a base size into a plurality of logical memory banks. A request to execute operations on the plurality of physical memory banks is received. The request to execute the operations comprises a request to interact with data having a sample width based on the base size. Responsive to receiving the request to execute the operations, the operations are executed on one or more logical memory banks of the plurality of logical memory banks via a memory crossbar shared across the plurality of logical memory banks. An amount of the data on which the operations are executed is a multiple of the sample width, and each logical memory bank has a size based on the base size and a multiplier value.

Inventors: Awasthi; Aditya (San Diego, CA), Yu; Tsung-Han (Lake Forest, CA), Li; Troy (Irvine, CA), Balasubramanian; Sundar Rajan (Groton, MA), Nayak; Ankita (Milpitas, CA), Kang; Leiter (Somerville, MA)

Applicant: QUALCOMM Incorporated (San Diego, CA)

Family ID: 1000008778808

Assignee: QUALCOMM Incorporated (San Diego, CA)

Appl. No.: 18/192033

Filed: March 29, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240330212 A1	Oct. 03, 2024

Publication Classification

Int. Cl.: G06F13/16 (20060101); G06F15/173 (20060101)

U.S. Cl.:

CPC G06F13/1657 (20130101); G06F13/1668 (20130101); G06F15/17375 (20130101);

Field of Classification Search

CPC: G06F (13/1657); G06F (13/1668); G06F (15/17375)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
12020035	12/2023	Parle	N/A	G06F 9/3887
12141082	12/2023	Minkin	N/A	G06F 12/0813
2013/0268715	12/2012	Fetterman	711/E12.078	G06F 9/3888
2013/0286706	12/2012	Perego	365/230.03	G06F 12/04
2022/0137866	12/2021	Kim	711/154	G11C 11/4076

OTHER PUBLICATIONS

International Search Report and Written Opinion—PCT/US2024/013573—ISA/EPO—May 31, 2024. cited by applicant

Aho E., et al., “Parallel Memory Architecture for Arbitrary Stride Accesses”, 2006 IEEE Design and Diagnostics of Electronic Circuits and Systems, Prague, Czech Republic, Apr. 18-21, 2006, 6 Pages. cited by applicant

Primary Examiner: Tsai; Henry

Assistant Examiner: Daley; Christopher A

Attorney, Agent or Firm: Patterson + Sheridan LLP

Background/Summary

INTRODUCTION

(1) Machine learning is generally the process of producing a trained model (e.g., an artificial neural network, a tree, or other structures), which represents a generalized fit to a set of training data. Applying the trained model to input data produces inferences, which may be used to gain insights into the input data. In some cases, applying the model to the input data is described as “running an inference” or “performing an inference” on the input data.

(2) To train a model and perform inferences on input data, various mathematical operations are performed using various mathematical processing components. For example, multiply-and-accumulate (MAC) units may be used to perform these operations to train a model and perform inferences on input data using the trained model. It should be noted, however, that MAC units may be used for various mathematical operations and are not so limited to use in mathematical operations related to training a model and performing inferences on input data. These mathematical operations are generally performed with respect to one or more matrices; for example, these

mathematical operations may include matrix-to-vector multiplication operations and matrix-to-matrix multiplication operations.

(3) Generally, the data stored in input matrices on which mathematical operations are performed and the data stored in output matrices generated by these mathematical operations are stored in a consistent format (e.g., in row major order, in which consecutive elements of a row of a matrix reside next to each other, or in column major order, in which consecutive elements of a column of a matrix reside next to each other). However, samples used in matrix operations (e.g., convolution operations in a machine learning model or the like) may be accessed in a specific order different from the order in which such data is stored in memory. Thus, to access data involved in matrix operations, various actions may be performed, such as sample packing or unpacking or data rearrangement, to allow for relevant samples to be obtained from or written to memory. These actions, however, impose various processing and power overheads and may increase the amount of time involved in performing various mathematical operations used in training a machine learning model and generating inferences using a machine learning model.

(4) Accordingly, what is needed are improved techniques for performing matrix operations more efficiently.

BRIEF SUMMARY

(5) Certain aspects of the present disclosure provide a method for efficiently accessing memory in a computing system. An example method includes organizing a plurality of physical memory banks having a base size into a plurality of logical memory banks. A request to execute operations on the plurality of physical memory banks is received. Generally, the request to execute the operations comprises a request to interact with data having a sample width based on the base size. Responsive to receiving the request to execute the operations, the operations are executed on one or more logical memory banks of the plurality of logical memory banks via a memory crossbar shared across the plurality of logical memory banks. Generally, an amount of the data on which the operations are executed is a multiple of the sample width, and each logical memory bank of the one or more logical memory banks has a size based on the base size and a multiplier value.

(6) Other aspects provide processing systems configured to perform the aforementioned methods as well as those described herein; non-transitory, computer-readable media comprising instructions that, when executed by one or more processors of a processing system, cause the processing system to perform the aforementioned methods as well as those described herein; a computer program product embodied on a computer-readable storage medium comprising code for performing the aforementioned methods as well as those further described herein; and a processing system comprising means for performing the aforementioned methods as well as those further described herein.

(7) The following description and the related drawings set forth in detail certain illustrative features of one or more aspects.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The appended figures depict only certain aspects of this disclosure and are therefore not to be considered limiting of the scope of this disclosure.

(2) FIG. 1 illustrates an example data access pattern for performing matrix operations in a fully connected layer of a machine learning model.

(3) FIG. 2 illustrates an example memory access architecture for accessing samples in memory according to a defined offset and stride length.

(4) FIG. 3 illustrates an example memory access architecture for accessing variable amounts of data at variable stride lengths using logical memory banks, according to aspects of the present

disclosure.

(5) FIG. 4 illustrates pseudocode for reading data from memory organized using logical memory banks based on a variable data width and variable stride length, according to aspects of the present disclosure.

(6) FIG. 5 illustrates pseudocode for writing data to memory organized using logical memory banks based on a variable data width and variable stride length, according to aspects of the present disclosure.

(7) FIG. 6 illustrates an example of memory access using variable data read widths and variable stride lengths via a shared memory crossbar, according to aspects of the present disclosure.

(8) FIG. 7 illustrates details of reading data from processor memory into buffers in which data is accessible via logical memory banks using variable sample width and variable stride length, according to aspects of the present disclosure.

(9) FIG. 8 illustrates example operations for memory organization and access for efficient matrix arithmetic, according to aspects of the present disclosure.

(10) FIG. 9 illustrates an example two-dimensional convolution performed using logical memory banks accessible via a shared memory crossbar, according to aspects of the present disclosure.

(11) FIG. 10 illustrates an example convolution with dilation and padding performed using logical memory banks accessible via a shared memory crossbar, according to aspects of the present disclosure.

(12) FIG. 11 illustrates an example recurrent neural network (RNN) cell with operations performed using logical memory banks accessible via a shared memory crossbar, according to aspects of the present disclosure.

(13) FIG. 12 depicts an example processing system configured to perform various aspects of the present disclosure.

(14) To facilitate understanding, identical reference numerals have been used, where possible, to designate identical elements that are common to the drawings. It is contemplated that elements and features of one aspect may be beneficially incorporated in other aspects without further recitation.

DETAILED DESCRIPTION

(15) Aspects of the present disclosure provide apparatus, methods, processing systems, and computer-readable mediums for organizing and accessing memory for efficient matrix operations (e.g., in training and inferencing using machine learning models).

(16) As discussed, many operations performed while training or inferencing using a machine learning model involve matrix operations, such as multiplication operations that result in the generation of a vector from a matrix or multiplication operations that result in the generation of a matrix from another matrix of data. These operations generally involve accessing samples of data from memory, which may not be organized in a manner conducive to accessing the samples of data on which these operations are to be performed. Thus, to allow for these samples in memory to be accessed, various operations to move or otherwise rearrange data in memory (also known as sample packing or unpacking) may be performed. These operations may use significant amounts of processing time, which may delay the commencement of operations performed on the samples accessed from memory. Further, moving or rearranging data in memory generally consumes some amount of power, which may impact the amount of energy that remains (e.g., in a battery or other energy storage device from which a processor draws power to operate) for other operations.

(17) Aspects of the present disclosure provide techniques for organizing and accessing data in memory that allow for efficient access to data samples in memory and accelerate the process of accessing data for use in various matrix operations, such as those involved in training a machine learning model or inferencing operations using a machine learning model. As discussed in further detail herein, physical memory may be organized into a plurality of logical memory banks having a base size and accessible via a memory crossbar shared across the plurality of logical memory banks. Incoming requests to interact with data may have a sample width that is based on the base

size, and the appropriate number of logical memory banks can be accessed to retrieve data from or write data to memory. By doing so, aspects of the present disclosure allow for an arbitrary number of contiguous samples of data in memory, with arbitrarily sized gaps (also known as a stride interval) between these samples, to be accessed in order to provide data to a processor for use in performing mathematical or other operations based on these samples. Because arbitrary numbers of contiguous samples of data can be accessed with arbitrarily sized stride intervals, aspects of the present disclosure may allow for data to be accessed without data sample packing or unpacking, which may reduce computational expense and power utilization involved in rearranging and moving data in memory to accommodate data access operations on arbitrary numbers of samples and arbitrary stride lengths.

(18) Introduction to Accessing Data in Memory

(19) FIG. 1 illustrates an example 100 of accessing data in memory (e.g., during machine learning operations, such as training a machine learning model or inferencing using a trained machine learning model) based on different stride lengths between data samples. As illustrated, example 100 illustrates the calculation of matrix Y **110** based on samples W **120**, X **130**, and B **140**. As illustrated, sample W **120** is a matrix with dimensions $M \times K$, sample X **130** is a matrix with dimensions $K \times N$, and sample B **140** is a matrix with dimensions $M \times N$.

(20) During a particular clock cycle, a processor generally accesses different portions of matrices including relevant samples of data. For example, during a given clock cycle, a processor can access a subsample **122** with dimensions $m \times k$, $m \in M$ and $k \in K$, from sample W **120**. Likewise, the processor can access a subsample **132** with dimensions $k \times N$ from sample X **130** and can access a subsample **142** with dimensions $m \times N$ from sample **140**. The resulting sum of subsample **142** and the product of the subsamples **122**, **132** from sample W **120** and from sample X **130** may be a matrix **112** with size $m \times N$.

(21) To access each subsample **122**, **132**, and **142**, some amount of memory may be read, and intervening memory need not be read. For example, to access the $m \times k$ subsample **122** from sample W **120**, the processor may read some elements from the third column of data, skip the remaining elements from the third column of data, then read some elements from the fourth column of data, and skip the remaining elements from the fourth column of data. Similarly, to access the $k \times N$ subsample **132** from sample X **130**, the processor may read some elements from each column of data and skip the remaining elements from each column of data. If the memory is organized in a column-major format, then retrieving the subsample **122** from sample W **120** may be represented as an access of m consecutive samples with a stride length of $M-m$, while retrieving the subsample **132** from sample X **130** may be represented as an access of k consecutive samples with a stride length of $K-k$ and an offset of 2.

(22) Data access patterns **150** and **160** illustrate examples of data access patterns in general matrix multiply (GEMM) operations in a fully connected layer of a neural network with rank- k updates (where the rank parameter corresponds to an amount of information flowing between layers of the neural network) without and with a beginning offset applied to the data access pattern, respectively. It should be recognized that the m and M parameters illustrated in data access pattern **150** may or may not be different from those illustrated in matrix Y **110**, and the k and K parameters illustrated in data access pattern **160** may be different from those illustrated in sample X **130**. As illustrated, data access pattern **150** illustrates accessing $m \in M$ samples from memory with no starting offset. With $m=8$ and $M=16$, it can be seen that within each sixteen-element sample (e.g., a first sample with indices 0 through 15, or a second sample with indices 16 through 31), the first eight elements are accessed (read from or written to), and the remaining eight elements are skipped. Similarly, data access pattern **160** illustrates accessing $k \in K$ samples from memory, but with a starting offset of 4 elements. With $k=4$, $K=12$, and the starting offset of 4, it can be seen that memory access does not start until element 4 (e.g., skips elements 0 through 3). Subsequently, for every block of 12 samples, a processor can access the first four samples and skip the remaining eight samples.

Example Arbitrary Stride Access in Memory

(23) To support accessing arbitrarily sized samples in memory and arbitrary stride lengths between accessed samples in memory, aspects of the present disclosure use a parallel memory architecture that allows a processor to request access to data in a number of parallel memory modules.

(24) FIG. 2 illustrates a memory access architecture **200** that allows for accessing arbitrarily sized samples of data in memory with arbitrary stride lengths between these samples of data. Memory access architecture **200** generally includes an address computation block **210** that generates a set of addresses identifying memory blocks **220** from which data is to be accessed (e.g., read from or written to). To generate the set of addresses, address computation block **210** uses three inputs: r , corresponding to a starting offset and identifying the initial location from which data is to be accessed, Stride.sub.a , corresponding to a stride length, and Stride.sub.s , corresponding to a skewing scheme defining how data is to be permuted for access.

(25) Read data permutation block **230** and write data permutation block **240** are configured, respectively, to read data from or write data to the addresses associated with memory blocks **220** generated by address computation block **210**. In some aspects, read data permutation block **230** and write data permutation block **240** can be controlled (e.g., activated or deactivated) using the $r\text{Control}$ and $w\text{Control}$ signals from address computation block **210**, respectively. If the $r\text{Control}$ is set to a value that activates the read data permutation block **240**, then read data permutation block **230** can retrieve data from N memory blocks 220.sub.0 through 220.sub.N-1 as N data bits rd.sub.0 through rd.sub.N-1 , respectively, and output these data bits to a processor for processing. If the write enable bit is set to a logical high, and $w\text{Control}$ activates the write data permutation block **240**, then N data bits wd.sub.0 through wd.sub.N-1 may be written to memory blocks 220.sub.0 through 220.sub.N-1 , respectively.

(26) Generally, memory access architecture **200** allows for memory addresses to be accessed according to a defined skewing scheme. Memory access architecture **200** allows for a data sample to be accessed, and then a subsequent data sample may be accessed after a hop over some amount of data from the previously accessed data sample. However, in many applications, such as training and inferencing using machine learning models, data samples of variable lengths (instead of a single data sample of a standardized length) may be processed in matrix operations, and variable stride lengths may be used to access the appropriate data samples. For example within a machine learning model, different layers may have differently sized inputs, outputs, and/or filters, which may involve the use of different memory access parameters, even though input and output matrices may be organized in a consistent format.

(27) To allow for efficient memory access that provides for accessing variable amounts of data in memory using variable stride lengths, aspects of the present disclosure provide techniques for partitioning or otherwise organizing memory into different logical memory banks that map to underlying physical memory banks. A single memory crossbar may be used to access one or more physical memory banks according to a number of memory samples over which data is to be accessed. To allow for arbitrary amounts of data (e.g., arbitrary sample widths) to be accessed, aspects of the present disclosure provide techniques to allow for a consecutive number of samples m to be accessed from arbitrarily sized sets of M data samples, so long as M is a multiple of m (i.e., so long as $M \bmod m = 0$).

(28) FIG. 3 illustrates an example memory architecture **300** in which a common memory crossbar allows for access to physical memory banks based on logical memory banks of varying sizes, according to aspects of the present disclosure. As illustrated, memory architecture **300** includes a plurality of physical memory banks **310A**, **310B**, . . . , **310C**, **310D** (collectively referred to herein as “physical memory banks **310**”) sharing a common crossbar. As illustrated, the plurality of physical memory banks **310** may allow for a base amount of memory to be accessed, and each of the logical memory banks may be accessed via one or more of the physical memory banks based on the size of the data being accessed. As illustrated, each memory bank may correspond to a 32-bit

storage unit which can output a 32-bit word as two 16-bit words, labeled out0 and out1. The two 16-bit words may be stored in different halves of the 32-bit storage unit (e.g., such that the 16 bits of out0 are stored in a bottom half of the 32-bit storage unit and the 16 bits of out1 are stored in a top half of the 32-bit storage unit, or vice versa). The common memory crossbar may be implemented by the plurality of physical memory banks **310** to allow for arbitrary sample widths to be accessed from memory. For example, as illustrated by memory architecture **300**, the common memory crossbar may include the plurality of physical memory banks **310**, allowing for multiples of a 16-bit word to be output or retrieved via the appropriate portions of each physical memory bank **310** (which, as discussed, supports two 16-bit words which in combination may be stored in different halves of a 32-bit storage unit).

(29) As discussed, to allow for arbitrary sample widths to be accessed from memory, logical memory banks may have a size that is defined as some multiple of a sample width, subject to the minimum size of a logical memory bank. For example, the size of a logical memory bank may be no smaller than the base width of a physical memory bank, and thus, even if the sample width m is smaller than the base width of the physical memory bank, the minimum amount of data returned from memory may be the amount of data stored in a single physical memory bank.

(30) In this example, each physical memory bank **310** is structured as a plurality of memory sub-banks that outputs a 16-bit data sample (e.g., each output is a sample with bit indices of 0 through 15). Because physical memory banks **310** are configured to output a 16-bit data sample, logical memory banks may be formed as multiple 16-bit samples. Assume, as illustrated in FIG. 3, that memory architecture **300** has a total memory capacity of 1024 bits. A first logical memory bank set **320** may thus include 64 16-bit memory banks, with each of logical memory banks **320A-320H** (and others not illustrated) corresponding to a physical memory bank **310**. A second logical memory bank set **330** may include half the number of memory banks as first logical memory bank set **320**, with each logical memory bank **330A-330H** (and others not illustrated) being composed of two logical memory banks in set **320**. A third logical memory bank set **340** may include half the number of memory banks as second logical memory bank set **330**, with each logical memory bank **340A-340H** (and others not illustrated) being composed of two logical memory banks in set **330** (and thus of four logical memory banks in set **320**). Finally, a fourth logical memory bank set **350** may include half the number of memory banks as third logical memory bank set **340**, with each logical memory bank **350A-350H** being composed of two logical memory banks in set **340** (and thus of four logical memory banks in set **330**, and ultimately, eight logical memory banks in set **320**).

(31) As illustrated, thus, memory architecture **300** allows for different amounts of memory to be accessed using a common memory crossbar. For example, when 32-bit words are selected, index 0 may correspond to indices 0 and 1 for the 16-bit words, index 1 may correspond to indices 2 and 3 for the 16-bit words, and so on. In this case, 16-bit word index 0 may correspond to the bottom half of the word at 32-bit word index 0, and 16-bit word index 1 may correspond to the top half of the word at 32-bit word index 0 (or vice-versa), and so on. By doing so, aspects of the present disclosure allow for memory to be arranged so that dynamic access patterns in terms of both sample width and stride length between samples can change with minimal, or at least reduced, memory reorganization. This support for dynamic access patterns may thus efficiently allow for dynamic sample shuffling and other operations used in training and inferencing operations by various layers of a machine learning model. Further, the common memory crossbar and various logical memory banks that can be supported using the common memory crossbar can allow for efficient memory access with different model dimensions, data bit widths, and underlying accelerator architectures.

(32) FIG. 4 illustrates example pseudocode **400** for reading data from memory organized using logical memory banks based on a variable data width and variable stride length, according to aspects of the present disclosure.

(33) As illustrated, pseudocode **400** begins with block **402** defining different data widths supported by the logical memory banks into which data is organized in memory. In this example, for a memory unit of 1024 bits, identifier B64 defines an operating mode with 64 logical memory banks with a capacity of 16 bits in each bank; identifier B32 defines an operating mode with 32 logical memory banks with a capacity of 32 bits in each bank; identifier B16 defines an operating mode with 16 logical memory banks with a capacity of 64 bits in each bank; and identifier B8 an operating mode with defines 8 logical memory banks with a capacity of 128 bits in each bank. It should be recognized, however, that other bank sizes may also be defined with other capacities, so long as the capacity of a bank is a multiple of a base size (which may correspond to the minimum width of a sample to be retrieved from memory).

(34) Pseudocode **400** may proceed to block **404**, in which the logical memory banks are configured to use a common crossbar. At block **404**, a read line adjustment and a read bank adjustment are defined for each data width defined in block **402**. These adjustments may be performed for each logical memory bank at the smallest size, since larger memory banks may be defined as multiples of smaller memory banks. Thus, as illustrated in pseudocode **400**, the read line adjustment and read bank adjustment for each logical memory bank size may be defined in terms of adjustments made so that the read line and read bank adjustments allow for data to be read starting from the correct start point and covering the appropriate number of logical memory banks (e.g., from the B64 mode, corresponding to the smallest-sized logical memory banks in memory).

(35) Finally, pseudocode **400** proceeds to block **406**, in which data is read from the common crossbar.

(36) FIG. 5 illustrates example pseudocode **500** for writing data to memory organized using logical memory banks based on a variable data width and variable stride length.

(37) As illustrated, pseudocode **500** begins with block **502** defining different data widths supported by the logical memory banks into which data is organized in memory. In this example, for a memory unit of 1024 bits, identifier MODE_64BK defines an operating mode with 64 logical memory banks with a capacity of 16 bits in each bank; identifier MODE_32BK defines an operating mode with 32 logical memory banks with a capacity of 32 bits in each bank; identifier MODE_16BK defines an operating mode with 16 logical memory banks with a capacity of 64 bits in each bank; and identifier MODE_8BK defines an operating mode with 8 logical memory banks with a capacity of 128 bits in each bank. It should be recognized, however, that other bank sizes may also be defined with other capacities, so long as the capacity of a bank is a multiple of a base size (which may correspond to the minimum width of a sample to be written to memory).

(38) Pseudocode **500** then proceeds to block **504**, in which data is written to memory based on the operating mode in which a write operation is to be performed. Block **504** generally identifies, based on an index adjustment and the mode in which operations are being performed, the locations in which data is to be written. As discussed, because larger blocks may be considered as multiple smaller blocks, and because the size of any memory bank is a multiple of the width of the smallest logical memory bank, write operations may be performed on any size of logical memory bank using a common memory crossbar that allows for access to memory based on the width of the smallest logical memory bank.

(39) FIG. 6 illustrates an example architecture **600** for memory access using variable data read widths and variable stride lengths via a shared memory crossbar, according to aspects of the present disclosure. Architecture **600** may, for example, be implemented in various compute architectures, such as neural processing units (NPU)s or other machine learning accelerators in which variable data read widths and stride lengths may be used to read data from and write data to memory. As illustrated, architecture **600** includes processor memory **610**, from which read access unit **620** accesses and provides data to an internal buffer (one or both of weight buffer **650** and input/output (I/O) buffer **660**) controlled by internal buffer write controller **630** and internal buffer read controller **670**.

(40) Generally, read access unit **620**, weight buffer **650**, and I/O buffer **660** can provide compute units **680** with data from one or both of the weight buffer **650** and the I/O buffer **660** based on variable amounts of data to be accessed (or sample width) and variable stride lengths between such data. The weights in weight buffer **650** may correspond to weights of a machine learning model loaded into memory for the compute units **680** to use in generating an inference result with respect to an input stored in I/O buffer **660**. Both the weights in weight buffer **650** and the input stored in I/O buffer **660** may be organized as matrices, portions of which may be multiplied into an output by compute units **680** and written back to processor memory **610** via write access unit **690**. To support accessing variable amounts of data with variable stride lengths between accessed data, the memory in which weight buffer **650** and I/O buffer **660** are located may be organized (e.g., partitioned) into a plurality of logical memory banks, which may be accessed in parallel using a common crossbar, as discussed above. These logical memory banks may have a size corresponding to the width of the underlying physical memory banks in which the data in weight buffer **650** and I/O buffer **660** are stored. Further, as discussed, one or more other layers of larger logical memory banks may be defined to allow for accessing larger amounts of data.

(41) During processing operations by compute units **680**, the results of matrix multiplication (or other matrix operations) generated for a subset of data from weight buffer **650** and a subset of data from I/O buffer **660** may be output and written back to I/O buffer **660** (or another buffer, not illustrated in FIG. 6). Thus, intermediate results of matrix operations generated by compute units **680** can remain buffered (e.g., in I/O buffer **660**) for access using the variable sample width and variable stride length memory access techniques discussed herein. When operations at compute units **680** are completed, the data in I/O buffer **660** may be output to processor memory **610** via write access unit **690**, which may be configured to write a matrix in a linear format (e.g., in row-major or column-major format) for storage and future access.

(42) FIG. 7 illustrates details of reading data from processor memory into buffers in which data is accessible via logical memory banks using the variable sample width and variable stride length techniques discussed herein, according to aspects of the present disclosure.

(43) Pipeline **700** illustrates a pipeline for reading static weight data into a weight buffer **650** and providing a subset of weight data from weight buffer **650** to compute units (e.g., compute units **680** illustrated in FIG. 6). Generally, two-dimensional weight matrix data may be stored in processor memory (e.g., processor memory **610** illustrated in FIG. 6) in linear form (e.g., in row-major or column-major format). Linear write block **710** generally converts the weight data from linear form as stored in processor memory to a form that is conducive to variable size and variable stride length access via logical memory banks and a common memory crossbar. After data is stored in weight buffer **650**, read access to subsets of data in weight buffer **650** may be performed by read data permute network **720**, which uses information about the size (or width) of a data sample requested by a compute unit and a stride length to identify logical memory banks from which data is to be retrieved. The identified logical memory banks may be read by stride read unit **730** and output to the compute unit(s) that use the subset of weight data in one or more matrix operations to generate an intermediate or final output, which may represent a result of an inference.

(44) Pipeline **705** illustrates a pipeline for reading input data and providing a subset of input data from I/O buffer **660** to compute units (e.g., compute units **680** illustrated in FIG. 6). As with the weight data discussed above, input data may also be stored in processor memory (e.g., processor memory **610** illustrated in FIG. 6) in linear form, and linear write block **710** converts the input data from linear form as stored in processor memory to a form that is conducive to variable size and variable stride length access via logical memory banks and a common memory crossbar. After data is stored in I/O buffer **660**, read access to subsets of data in I/O buffer **660** may be performed by read data permute network **720**, which uses information about the size (or width) of a data sample requested by a compute unit and a stride length to identify logical memory banks from which data is to be retrieved. The identified logical memory banks may be read by stride read unit **730** and

output to the compute unit(s) that use the subset of input data in one or more matrix operations to generate an intermediate or final output, which may represent a result of an inference.

(45) In some aspects, the results generated by the compute units based on the subset of input data output to the compute units by stride read unit **730** may be used in subsequent operations performed by these compute units. For example, the results may represent the output of one layer of a machine learning model which is used as input into another layer of the machine learning model. To allow for continued efficiencies in accessing such data, a write data permute network **740** can identify the logical memory blocks in I/O buffer **660** to which data is to be written and the stride length between different locations at which data is to be written. Stride write unit **750** may then write the results generated by the compute units to I/O buffer **660** at the identified addresses, and the read/write process discussed herein may continue while machine learning operations with respect to an input are being performed.

(46) In some aspects, whether data is written to I/O buffer **660** from processor memory via linear write block **710** or from compute units via stride write unit **750** may be controlled by multiplexer **715**. For example, multiplexer **715** may be set to allow I/O buffer **660** to be populated with input data from processor memory via linear write block **710** initially. After the input data is written from processor memory to I/O buffer **660**, multiplexer **715** may be set to allow I/O buffer **660** to be populated or updated with data from the compute units via stride write unit **750**. The multiplexer **715** may generally be set to populate I/O buffer **660** with data from processor memory or from the compute units on an as-demanded basis.

(47) FIG. **8** illustrates example operations **800** for memory organization and access for efficient matrix arithmetic, according to aspects of the present disclosure. Operations **800** may be performed, for example, by a computing device on which matrix arithmetic operations are performed for various tasks, such as processing system **1200** shown in FIG. **12**. This computing device may include a user equipment (UE) or other edge device, a server, cloud computing resources, or another device on which matrix math operations are performed (e.g., in training and/or inferencing using machine learning models).

(48) As illustrated, operations **800** may begin at block **810**, with organizing a plurality of physical memory banks having a base size into a plurality of logical memory banks.

(49) In some aspects, the plurality of logical memory banks may be structured as an architecture of multiple sets of logical memory banks. A first set of logical memory banks may have a size equal to the base size (as it may not be practicable for logical memory banks to have a size smaller than the size of the underlying physical memory banks). A second set of logical memory banks may have a size double ($2\times$) that of the base size (or some other multiple of the base size). Additional sets of logical memory banks may be defined with increasing sizes (and thus decreasing numbers of memory banks for a given amount of memory); for example, a third set of logical memory banks may have a size quadruple ($4\times$) that of the base size; a fourth set of logical memory banks may have a size octuple ($8\times$) that of the base size; and so on. More generally, the logical banks may provide access to up to M contiguous samples of data, where M is a multiple of a minimum sample width m .

(50) In some aspects, the plurality of physical memory banks may be organized into a plurality of logical memory banks a priori and accessed via a logical bank mode. In some aspects, the organizing (e.g., partitioning) of physical memory banks to logical memory banks may be performed dynamically and in real time in response to determining that such an organization would be useful in accelerating memory access for matrix operations.

(51) At block **820**, operations **800** proceed with receiving a request to execute operations on the plurality of physical memory banks, wherein the request to execute the operations comprises a request to interact with data having a sample width based on the base size.

(52) In some aspects, the request to execute operations includes a request to read data from the physical memory banks. In such a case, to execute the operations on the plurality of physical

memory banks, data is retrieved from at least one logical memory bank of the second set of logical memory banks via the shared memory crossbar. To retrieve the data from the at least one logical memory bank via the shared crossbar, data is retrieved from logical memory banks of the first set of logical memory banks that form the at least one logical memory bank of the second set of logical memory banks. The retrieved data may be output to one or more compute units configured to process the retrieved data using a machine learning model.

(53) In some aspects, the request to execute operations comprises a request to write data to the plurality of physical memory banks. In such a case, to execute the operations on the plurality of physical memory banks, the data is written to at least one logical memory bank of the second set of logical memory banks via the shared memory crossbar. To write the data to the at least one logical memory bank of the second set of logical memory banks, the data is written to logical memory banks of the first set of logical memory banks that form the at least one logical memory bank of the second set of logical memory banks.

(54) At block **830**, responsive to receiving the request to execute the operations, operations **800** proceed with executing the operations on one or more logical memory banks of the plurality of logical memory banks via a memory crossbar shared across the plurality of logical memory banks. An amount of the data on which the operations are executed is a multiple of the sample width, and each logical memory bank of the one or more logical memory banks has a size based on the base size and a multiplier value.

(55) In some aspects, the request to execute the operations on the plurality of physical memory banks is received from one or more compute units configured to process the written data using a machine learning model.

(56) In some aspects, the request to execute the operations on the plurality of physical memory banks is based on data from the one or more compute units or data from a memory associated with a processor, selected via a multiplexer.

(57) In some aspects, operations **800** further include receiving a request to execute other operations on the plurality of physical memory banks. The request to execute the other operations may include a request to interact with other data located in a non-contiguously located physical memory bank relative to the plurality of physical memory banks on which the operations are executed. Responsive to receiving the request to execute the other operations, the other operations are executed on the one or more logical memory banks of the plurality of logical memory banks via the shared memory crossbar.

Example Machine Learning Model Operations

(58) FIG. **9** illustrates an example two-dimensional convolution performed using logical memory banks accessible via a shared memory crossbar, according to aspects of the present disclosure. A convolution operation is a common operation performed by layers of a machine learning model during training and inferencing operations.

(59) As illustrated, convolution **900** generates an output **910** based on multiple passes **920.sub.1-920.sub.n**, multiplied by an input **930** (or at least a portion of the input **930**). Each of these passes **920.sub.1-920.sub.n** corresponds to the application of a different kernel in the neural network to input **930**, allowing for the generation of different two-dimensional output planes in output **910**. To allow for output **910** to be generated using convolutions performed on data stored in logical memory banks accessible via a shared memory crossbar, a two-dimensional convolution may be decomposed into multiple general matrix multiply (GEMM) operations corresponding to passes **920.sub.1-920.sub.n**. Each pass, in some aspects, may generate data based on the application of a different kernel to of a portion of input **930**, represented by the shaded portion of input **930**, in order to generate values for at least the shaded portions of output **910**. Each pass **920** may compute m output channels and n output pixels and may be formulated as an unrolling of each kernel in the neural network as a row in an unrolled weight matrix **940** and unrolling each activation as a column of an unrolled activation matrix **950**, where activation duplication can be achieved by reading the

same data from memory in different processing cycles. For example, as illustrated, a given pass representing a portion of a two-dimensional convolution may compute $m=8$ output channels and $n=4$ output pixels. The $m=8$ output channels may be generated based on the shaded rows of weight matrix **940**, and the $n=4$ output pixels may be generated based on the shaded columns of the activation matrix **950**. The number of GEMM operations into which a two-dimensional convolution may be decomposed may be calculated by the expression:

$$(60) \text{ ceiling}(\frac{O_{c'}}{m}) * \text{ ceiling}(\frac{O_{x'}}{sx}) * \text{ ceiling}(\frac{O_{y'}}{sy})$$

where $O_{c'}$ represents the number of channels in an output, $O_{x'}$ represents the number of pixels in an output on the x axis, sx represents a stride length over the x axis, $O_{y'}$ represents the number of pixels in an output on the y axis, and sy represents a stride length over the y axis.

(61) FIG. **10** illustrates an example convolution **1000** with dilation and padding performed using logical memory banks accessible via a shared memory crossbar, according to aspects of the present disclosure.

(62) As illustrated, the output **1010** of convolution **1000** may be based on a product of weight matrices **1020** and input activation matrices **1030**. Input activation matrices **1030** may represent a dilated input in which only portions of an input are used in convolution operations. For example, portions **1032**, **1034** of input activation matrices **1030** may be omitted from convolution operations, which allows for convolution **1000** to be performed over a larger area of an input activation matrix **1030**. In an image processing system, for example, a dilated input may comprise a portion of an image over which certain pixels are disregarded during a convolution operation, based on an assumption that skipping or otherwise disregarding these pixels during convolution operations will not result in significantly degraded inference performance because the skipped pixels generally include information that is similarly captured in the pixels that are used during convolution **1000**. The memory access techniques described herein may support the use of dilation and also of padding an input activation matrix with “dummy” bits **1036**, which may be bits not originally included in the input activation matrix **1030** but used to pad a matrix into a sufficient size for inference operations including convolution **1000**.

(63) FIG. **11** illustrates an example recurrent neural network (RNN) cell with operations **1100** performed using logical memory banks accessible via a shared memory crossbar, according to aspects of the present disclosure. RNN cells are generally building blocks for various types of recurrent networks used in applications such as natural language processing, time-series prediction, voice recognition, speech recognition, and the like. These recurrent networks may include gated recurrent units (GRUs), long short-term memory (LSTM) networks, and the like.

(64) As illustrated, operations **1100** in an RNN cell may be decomposed into an output **1110** that is the sum of two different GEMM passes and a vector addition operation performed prior to activation. A first GEMM pass may be, for example, represented as the product of weight matrix **1120** and a hidden layer output **1130**. A second GEMM pass may be, for example, represented as the product of weight matrix **1150** and input **1160**. The outputs of the first GEMM pass and the second GEMM pass, which may both be matrices, may be added together prior to applying activations on the output **1110**. Further, the outputs of the first GEMM pass and the second GEMM pass may be modified by bias matrices **1140** and **1170**, respectively.

(65) In some aspects, the inputs into a vector addition operation may be stored in a single memory. In such a case, the memory architecture described herein may allow for simultaneous access to both inputs, which may further accelerate the process of generating an inference using machine learning models operating on such data.

Example Processing Systems for Memory Organization and Access for Efficient Matrix Arithmetic

(66) FIG. **12** depicts an example processing system **1200** for memory organization and access for efficient matrix arithmetic, such as described herein for example with respect to FIG. **8**.

(67) Processing system **1200** includes a central processing unit (CPU) **1202**, which in some

examples may be a multi-core CPU. Instructions executed at the CPU **1202** may be loaded, for example, from a program memory associated with the CPU **1202** or may be loaded from a memory partition (e.g., of memory **1224**).

(68) Processing system **1200** also includes additional processing components tailored to specific functions, such as a graphics processing unit (GPU) **1204**, a digital signal processor (DSP) **1206**, a neural processing unit (NPU) **1208**, and a connectivity component **1212**.

(69) An NPU, such as NPU **1208**, is generally a specialized circuit configured for implementing control and arithmetic logic for executing machine learning algorithms, such as algorithms for processing artificial neural networks (ANNs), deep neural networks (DNNs), random forests (RFs), and the like. An NPU may sometimes alternatively be referred to as a neural signal processor (NSP), tensor processing unit (TPU), neural network processor (NNP), intelligence processing unit (IPU), vision processing unit (VPU), or graph processing unit.

(70) NPUs, such as NPU **1208**, are configured to accelerate the performance of common machine learning tasks, such as image classification, machine translation, object detection, and various other predictive models. In some examples, a plurality of NPUs may be instantiated on a single chip, such as a system on a chip (SoC), while in other examples such NPUs may be part of a dedicated neural-network accelerator.

(71) NPUs may be optimized for training or inference, or in some cases configured to balance performance between both. For NPUs that are capable of performing both training and inference, the two tasks may still generally be performed independently.

(72) NPUs designed to accelerate training are generally configured to accelerate the optimization of new models, which is a highly compute-intensive operation that involves inputting an existing dataset (often labeled or tagged), iterating over the dataset, and then adjusting model parameters, such as weights and biases, in order to improve model performance. Generally, optimizing based on a wrong prediction involves propagating back through the layers of the model and determining gradients to reduce the prediction error.

(73) NPUs designed to accelerate inference are generally configured to operate on complete models. Such NPUs may thus be configured to input a new piece of data and rapidly process this new piece through an already trained model to generate a model output (e.g., an inference).

(74) In one implementation, NPU **1208** is a part of one or more of CPU **1202**, GPU **1204**, and/or DSP **1206**. These may be located on a UE or another computing device.

(75) In some examples, connectivity component **1212** may include subcomponents, for example, for third generation (3G) connectivity, fourth generation (4G) connectivity (e.g., LTE), fifth generation (5G) connectivity (e.g., NR), Wi-Fi connectivity, Bluetooth connectivity, and other wireless data transmission standards. Connectivity component **1212** may be further coupled to one or more antennas **1214**.

(76) In some examples, one or more of the processors of processing system **1200** may be based on an ARM or RISC-V instruction set.

(77) Processing system **1200** also includes memory **1224**, which is representative of one or more static and/or dynamic memories, such as a dynamic random access memory, a flash-based static memory, and the like. In this example, memory **1224** includes computer-executable components, which may be executed by one or more of the aforementioned processors of processing system **1200**.

(78) In particular, in this example, memory **1224** includes memory bank organizing component **1224A**, request receiving component **1224B**, and operation executing component **1224C**. The depicted components, and others not depicted, may be configured to perform various aspects of the methods described herein.

(79) Generally, processing system **1200** and/or components thereof may be configured to perform the methods described herein.

EXAMPLE CLAUSES

(80) Implementation examples are described in the following numbered clauses: Clause 1: A computer-implemented method, comprising: organizing a plurality of physical memory banks having a base size into a plurality of logical memory banks; receiving a request to execute operations on the plurality of physical memory banks, wherein the request to execute the operations comprises a request to interact with data having a sample width based on the base size; and responsive to receiving the request to execute the operations, executing the operations on one or more logical memory banks of the plurality of logical memory banks via a memory crossbar shared across the plurality of logical memory banks, wherein an amount of the data on which the operations are executed is a multiple of the sample width, and wherein each logical memory bank of the one or more logical memory banks has a size based on the base size and a multiplier value. Clause 2: The method of Clause 1, wherein the plurality of logical memory banks comprises: a first set of logical memory banks having a size equal to the base size; and a second set of logical memory banks having a size double that of the base size. Clause 3: The method of Clause 2, wherein: the request to execute operations further comprises a request to read data from the plurality of physical memory banks, executing the operations comprises retrieving the data from at least one logical memory bank of the second set of logical memory banks via the shared memory crossbar, and retrieving the data from the at least one logical memory bank of the second set of logical memory banks comprises retrieving the data from logical memory banks of the first set of logical memory banks that form the at least one logical memory bank of the second set of logical memory banks. Clause 4: The method of Clause 3, further comprising outputting the retrieved data to one or more compute units configured to process the retrieved data using a machine learning model. Clause 5: The method of Clause 2, wherein: the request to execute operations comprises a request to write data to the plurality of physical memory banks, executing the operations comprises writing the data to at least one logical memory bank of the second set of logical memory banks via the shared memory crossbar, and writing the data to the at least one logical memory bank of the second set of logical memory banks comprises writing the data to logical memory banks of the first set of logical memory banks that form the at least one logical memory bank of the second set of logical memory banks. Clause 6: The method of Clause 5, wherein the request to execute the operations on the plurality of physical memory banks is received from one or more compute units configured to process the written data using a machine learning model. Clause 7: The method of Clause 6, wherein the request to execute the operations on the plurality of physical memory banks is based on data from the one or more compute units or data from a memory associated with a processor, selected via a multiplexer. Clause 8: The method of any of Clauses 1 through 7, further comprising: receiving a request to execute other operations on the plurality of physical memory banks, wherein the request to execute the other operations comprises a request to interact with other data located in a non-contiguously located physical memory bank relative to the plurality of physical memory banks on which the operations are executed; and responsive to receiving the request to execute the other operations, executing the other operations on the one or more logical memory banks of the plurality of logical memory banks via the shared memory crossbar. Clause 9: A system, comprising: a memory having executable instructions stored thereon; and a processor configured to execute the executable instructions in order to cause the system to perform the operations of any of Clauses 1 through 8. Clause 10: A system, comprising means for performing the operations of any of Clauses 1 through 8. Clause 11: A computer-readable medium having instructions stored thereon which, when executed by a processor, perform the operations of any of Clauses 1 through 8.

Additional Considerations

(81) The preceding description is provided to enable any person skilled in the art to practice the various aspects described herein. The examples discussed herein are not limiting of the scope, applicability, or aspects set forth in the claims. Various modifications to these aspects will be readily apparent to those skilled in the art, and the generic principles defined herein may be applied

to other aspects. For example, changes may be made in the function and arrangement of elements discussed without departing from the scope of the disclosure. Various examples may omit, substitute, or add various procedures or components as appropriate. For instance, the methods described may be performed in an order different from that described, and various steps may be added, omitted, or combined. Also, features described with respect to some examples may be combined in some other examples. For example, an apparatus may be implemented or a method may be practiced using any number of the aspects set forth herein. In addition, the scope of the disclosure is intended to cover such an apparatus or method that is practiced using other structure, functionality, or structure and functionality in addition to, or other than, the various aspects of the disclosure set forth herein. It should be understood that any aspect of the disclosure disclosed herein may be embodied by one or more elements of a claim.

(82) As used herein, the word “exemplary” means “serving as an example, instance, or illustration.” Any aspect described herein as “exemplary” is not necessarily to be construed as preferred or advantageous over other aspects.

(83) As used herein, a phrase referring to “at least one of” a list of items refers to any combination of those items, including single members. As an example, “at least one of: a, b, or c” is intended to cover a, b, c, a-b, a-c, b-c, and a-b-c, as well as any combination with multiples of the same element (e.g., a-a, a-a-a, a-a-b, a-a-c, a-b-b, a-c-c, b-b, b-b-b, b-b-c, c-c, and c-c-c or any other ordering of a, b, and c).

(84) As used herein, the term “determining” encompasses a wide variety of actions. For example, “determining” may include calculating, computing, processing, deriving, investigating, looking up (e.g., looking up in a table, a database or another data structure), ascertaining, and the like. Also, “determining” may include receiving (e.g., receiving information), accessing (e.g., accessing data in a memory), and the like. Also, “determining” may include resolving, selecting, choosing, establishing, and the like.

(85) The methods disclosed herein comprise one or more steps or actions for achieving the methods. The method steps and/or actions may be interchanged with one another without departing from the scope of the claims. In other words, unless a specific order of steps or actions is specified, the order and/or use of specific steps and/or actions may be modified without departing from the scope of the claims. Further, the various operations of methods described above may be performed by any suitable means capable of performing the corresponding functions. The means may include various hardware and/or software component(s) and/or module(s), including, but not limited to a circuit, an application specific integrated circuit (ASIC), or processor. Generally, where there are operations illustrated in figures, those operations may have corresponding counterpart means-plus-function components with similar numbering.

(86) The following claims are not intended to be limited to the aspects shown herein, but are to be accorded the full scope consistent with the language of the claims. Within a claim, reference to an element in the singular is not intended to mean “one and only one” unless specifically so stated, but rather “one or more.” Unless specifically stated otherwise, the term “some” refers to one or more. No claim element is to be construed under the provisions of 35 U.S.C. § 112(f) unless the element is expressly recited using the phrase “means for” or, in the case of a method claim, the element is recited using the phrase “step for.” All structural and functional equivalents to the elements of the various aspects described throughout this disclosure that are known or later come to be known to those of ordinary skill in the art are expressly incorporated herein by reference and are intended to be encompassed by the claims. Moreover, nothing disclosed herein is intended to be dedicated to the public regardless of whether such disclosure is explicitly recited in the claims.

Claims

1. A processor-implemented method, comprising: organizing a plurality of physical memory banks having a base size into a plurality of logical memory banks that comprise a first set of logical memory banks having a size equal to the base size and a second set of logical memory banks having a size double that of the base size; receiving a request to execute operations on the plurality of physical memory banks, wherein the request to execute the operations comprises a request to interact with data having a sample width based on the base size and a request to read the data from the plurality of physical memory banks; and responsive to receiving the request to execute the operations, executing the operations on one or more logical memory banks of the plurality of logical memory banks via a memory crossbar shared across the plurality of logical memory banks, wherein an amount of the data on which the operations are executed is a multiple of the sample width, wherein each logical memory bank of the one or more logical memory banks has a size based on the base size and a multiplier value, wherein executing the operations comprises retrieving the data from at least one first logical memory bank of the second set of logical memory banks via the shared memory crossbar, and wherein retrieving the data from the at least one first logical memory bank of the second set of logical memory banks comprises retrieving the data from logical memory banks of the first set of logical memory banks that form the at least one first logical memory bank of the second set of logical memory banks.
2. The method of claim 1, further comprising outputting the retrieved data to one or more compute units configured to process the retrieved data using a machine learning model.
3. The method of claim 1, wherein: the request to execute operations further comprises a request to write the data to the plurality of physical memory banks, executing the operations further comprises writing the data to at least one second logical memory bank of the second set of logical memory banks via the shared memory crossbar, and writing the data to the at least one second logical memory bank of the second set of logical memory banks comprises writing the data to logical memory banks of the first set of logical memory banks that form the at least one second logical memory bank of the second set of logical memory banks.
4. The method of claim 3, wherein the request to execute the operations on the plurality of physical memory banks is received from one or more compute units configured to process the written data using a machine learning model.
5. The method of claim 4, wherein the request to execute the operations on the plurality of physical memory banks is based on data from the one or more compute units or data from a memory associated with a processor, selected via a multiplexer.
6. The method of claim 1, further comprising: receiving a request to execute other operations on the plurality of physical memory banks, wherein the request to execute the other operations comprises a request to interact with other data located in a non-contiguously located physical memory bank relative to the plurality of physical memory banks on which the operations are executed; and responsive to receiving the request to execute the other operations, executing the other operations on the one or more logical memory banks of the plurality of logical memory banks via the shared memory crossbar.
7. A system, comprising: memory having executable instructions stored thereon; and one or more processors configured, individually or collectively, to execute the executable instructions in order to cause the system to: organize a plurality of physical memory banks having a base size into a plurality of logical memory banks that comprise a first set of logical memory banks having a size equal to the base size and a second set of logical memory banks having a size double that of the base size; receive a request to execute operations on the plurality of physical memory banks, wherein the request to execute the operations comprises a request to interact with data having a sample width based on the base size and a request to read the data from the plurality of physical memory banks; and responsive to receiving the request to execute the operations, execute the operations on one or more logical memory banks of the plurality of logical memory banks via a

memory crossbar shared across the plurality of logical memory banks, wherein an amount of the data on which the operations are executed is a multiple of the sample width, wherein each logical memory bank of the one or more logical memory banks has a size based on the base size and a multiplier value, wherein in order to execute the operations, the one or more processors are configured, individually or collectively, to cause the system to retrieve the data from at least one first logical memory bank of the second set of logical memory banks via the shared memory crossbar, and wherein in order to retrieve the data from the at least one first logical memory bank of the second set of logical memory banks, the one or more processors are configured, individually or collectively, to cause the system to retrieve the data from logical memory banks of the first set of logical memory banks that form the at least one first logical memory bank of the second set of logical memory banks.

8. The system of claim 7, wherein the one or more processors are further configured, individually or collectively, to cause the system to output the retrieved data to one or more compute units configured to process the retrieved data using a machine learning model.

9. The system of claim 7, wherein: the request to execute operations further comprises a request to write the data to the plurality of physical memory banks, in order to execute the operations, the one or more processors are further configured, individually or collectively, to cause the system to write the data to at least one second logical memory bank of the second set of logical memory banks via the shared memory crossbar, and in order to write the data to the at least one second logical memory bank of the second set of logical memory banks, the one or more processors are further configured, individually or collectively, to write the data to logical memory banks of the first set of logical memory banks that form the at least one second logical memory bank of the second set of logical memory banks.

10. The system of claim 9, wherein the request to execute the operations on the plurality of physical memory banks is received from one or more compute units configured to process the written data using a machine learning model.

11. The system of claim 10, wherein the request to execute the operations on the plurality of physical memory banks is based on data from the one or more compute units or data from another memory associated with the one or more processors, selected via a multiplexer.

12. The system of claim 7, wherein the one or more processors are further configured, individually or collectively, to cause the system to: receive a request to execute other operations on the plurality of physical memory banks, wherein the request to execute the other operations comprises a request to interact with other data located in a non-contiguously located physical memory bank relative to the plurality of physical memory banks on which the operations are executed; and responsive to receiving the request to execute the other operations, execute the other operations on the one or more logical memory banks of the plurality of logical memory banks via the shared memory crossbar.

13. A system, comprising: means for organizing a plurality of physical memory banks having a base size into a plurality of logical memory banks that comprise a first set of logical memory banks having a size equal to the base size and a second set of logical memory banks having a size double that of the base size; means for receiving a request to execute operations on the plurality of physical memory banks, wherein the request to execute the operations comprises a request to interact with data having a sample width based on the base size and a request to read the data from the plurality of physical memory banks; and means for executing, responsive to receiving the request to execute the operations, the operations on one or more logical memory banks of the plurality of logical memory banks via a memory crossbar shared across the plurality of logical memory banks, wherein an amount of the data on which the operations are executed is a multiple of the sample width, wherein each logical memory bank of the one or more logical memory banks has a size based on the base size and a multiplier value, wherein the means for executing the operations comprises means for retrieving the data from at least one first logical memory bank of the second

set of logical memory banks via the shared memory crossbar, and wherein the means for retrieving the data from the at least one first logical memory bank of the second set of logical memory banks comprises means for retrieving the data from logical memory banks of the first set of logical memory banks that form the at least one first logical memory bank of the second set of logical memory banks.

14. The system of claim 13, further comprising means for outputting the retrieved data to one or more compute units configured to process the retrieved data using a machine learning model.

15. The system of claim 13, wherein: the request to execute operations further comprises a request to write the data to the plurality of physical memory banks, the means for executing the operations further comprises means for writing the data to at least one second logical memory bank of the second set of logical memory banks via the shared memory crossbar, and the means for writing the data to the at least one second logical memory bank of the second set of logical memory banks further comprises means for writing the data to logical memory banks of the first set of logical memory banks that form the at least one second logical memory bank of the second set of logical memory banks.

16. The system of claim 15, wherein the means for receiving the request to execute the operations on the plurality of physical memory banks comprises means for receiving the request from one or more compute units configured to process the written data using a machine learning model.

17. The system of claim 16, wherein the request to execute the operations on the plurality of physical memory banks is based on data from the one or more compute units or data from memory associated with one or more processors, selected via a multiplexer.

18. The system of claim 13, further comprising: means for receiving a request to execute other operations on the plurality of physical memory banks, wherein the request to execute the other operations comprises a request to interact with other data located in a non-contiguously located physical memory bank relative to the plurality of physical memory banks on which the operations are executed; and means for executing, responsive to receiving the request to execute the other operations, the other operations on the one or more logical memory banks of the plurality of logical memory banks via the shared memory crossbar.

19. A non-transitory computer-readable medium having instructions stored thereon which, when executed by one or more processors individually or collectively, perform an operation comprising: organizing a plurality of physical memory banks having a base size into a plurality of logical memory banks that comprise a first set of logical memory banks having a size equal to the base size and a second set of logical memory banks having a size double that of the base size; receiving a request to execute operations on the plurality of physical memory banks, wherein the request to execute the operations comprises a request to interact with data having a sample width based on the base size and a request to read the data from the plurality of physical memory banks; and responsive to receiving the request to execute the operations, executing the operations on one or more logical memory banks of the plurality of logical memory banks via a memory crossbar shared across the plurality of logical memory banks, wherein an amount of the data on which the operations are executed is a multiple of the sample width, wherein each logical memory bank of the one or more logical memory banks has a size based on the base size and a multiplier value, wherein executing the operations comprises retrieving the data from at least one first logical memory bank of the second set of logical memory banks via the shared memory crossbar, and wherein retrieving the data from the at least one first logical memory bank of the second set of logical memory banks comprises retrieving the data from logical memory banks of the first set of logical memory banks that form the at least one first logical memory bank of the second set of logical memory banks.
