

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12399886
Kind Code	B2
Date of Patent	August 26, 2025
Inventor(s)	Chandrashekhhar; Ganesan et al.

Parsing logical network definition for different sites

Abstract

Some embodiments provide a method for distributing a service rule that is to be enforced across a first set of sites and that is defined by reference to a group identifier that identifies a group of machines. The method distributes the service rule to each site in the first set of sites. The method identifies at least one site in the first set of sites that is not in a second set of sites that has already received a definition of the group. The method distributes the group definition to each identified site in the first set of sites that has not already received the definition of the group.

Inventors: Chandrashekhhar; Ganesan (Campbell, CA), Margarian; Pavlush (Dublin, CA), Palavalli; Amarnath (Fremont, CA), Dorr; Josh (San Jose, CA)

Applicant: VMware LLC (Palo Alto, CA)

Family ID: 1000008778241

Assignee: VMware LLC (Palo Alto, CA)

Appl. No.: 18/243435

Filed: September 07, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20230421445 A1	Dec. 28, 2023

Foreign Application Priority Data

IN	202041015134	Apr. 06, 2020
----	--------------	---------------

Related U.S. Application Data

continuation parent-doc US 16906955 20200619 US 11799726 child-doc US 18243435

Publication Classification

Int. Cl.: **H04L41/0893** (20220101); **G06F9/54** (20060101); **G06F11/34** (20060101); **G06F16/23** (20190101); **G06F16/25** (20190101); **G06F16/27** (20190101); **H04L9/40** (20220101); **H04L12/66** (20060101); **H04L41/046** (20220101); **H04L41/0654** (20220101); **H04L41/0686** (20220101); **H04L41/0803** (20220101); **H04L41/0894** (20220101); **H04L41/0895** (20220101); **H04L41/12** (20220101); **H04L41/122** (20220101); **H04L41/22** (20220101); **H04L41/40** (20220101); **H04L43/06** (20220101); **H04L43/0811** (20220101); **H04L61/103** (20220101); **H04L67/1095** (20220101); **H04L67/14** (20220101); **H04L69/22** (20220101)

U.S. Cl.:

CPC **G06F16/2365** (20190101); **G06F9/546** (20130101); **G06F9/547** (20130101); **G06F11/3409** (20130101); **G06F16/2358** (20190101); **G06F16/256** (20190101); **G06F16/27** (20190101); **H04L12/66** (20130101); **H04L41/046** (20130101); **H04L41/0654** (20130101); **H04L41/0686** (20130101); **H04L41/0803** (20130101); **H04L41/0893** (20130101); **H04L41/0894** (20220501); **H04L41/0895** (20220501); **H04L41/12** (20130101); **H04L41/122** (20220501); **H04L41/40** (20220501); **H04L43/06** (20130101); **H04L43/0811** (20130101); **H04L61/103** (20130101); **H04L63/0263** (20130101); **H04L63/104** (20130101); **H04L67/14** (20130101); **H04L69/22** (20130101); **H04L41/22** (20130101); **H04L67/1095** (20130101)

Field of Classification Search

CPC: H04L (41/40); H04L (61/103); H04L (41/0654); H04L (67/14); H04L (41/0893); H04L (63/0263); G06F (16/2365)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5842043	12/1997	Nishimura	N/A	N/A
6219699	12/2000	McCloghrie et al.	N/A	N/A
6347374	12/2001	Drake et al.	N/A	N/A
6615230	12/2002	Nishimura	N/A	N/A
7502884	12/2008	Shah et al.	N/A	N/A
7539745	12/2008	Wang et al.	N/A	N/A
7802000	12/2009	Huang et al.	N/A	N/A
7965699	12/2010	Accardi et al.	N/A	N/A
8443363	12/2012	Brennan et al.	N/A	N/A
8479275	12/2012	Naseh	N/A	N/A
8566900	12/2012	Bharali	709/224	H04L 63/107
8611351	12/2012	Gooch et al.	N/A	N/A
8625616	12/2013	Vobbilisetty et al.	N/A	N/A
8660129	12/2013	Brendel et al.	N/A	N/A
8707417	12/2013	Liang et al.	N/A	N/A
9069599	12/2014	Martinez et al.	N/A	N/A
9215213	12/2014	Bansal et al.	N/A	N/A

9294270	12/2015	Wainner et al.	N/A	N/A
9311122	12/2015	Guay et al.	N/A	N/A
9330161	12/2015	D'Amato et al.	N/A	N/A
9432215	12/2015	Stabile et al.	N/A	N/A
9602312	12/2016	Koponen et al.	N/A	N/A
9672054	12/2016	Gupta et al.	N/A	N/A
9672060	12/2016	Behere et al.	N/A	N/A
9825851	12/2016	Agarwal et al.	N/A	N/A
9847938	12/2016	Chanda et al.	N/A	N/A
9876711	12/2017	Chu et al.	N/A	N/A
9906560	12/2017	Jain et al.	N/A	N/A
9912616	12/2017	Shen et al.	N/A	N/A
9923811	12/2017	Agarwal et al.	N/A	N/A
9977688	12/2017	Nipane et al.	N/A	N/A
10091028	12/2017	Koponen et al.	N/A	N/A
10110417	12/2017	Hankins et al.	N/A	N/A
10120668	12/2017	Palavalli et al.	N/A	N/A
10129142	12/2017	Goliya et al.	N/A	N/A
10135675	12/2017	Yu et al.	N/A	N/A
10142127	12/2017	Cherian et al.	N/A	N/A
10162656	12/2017	Palavalli et al.	N/A	N/A
10187302	12/2018	Chu et al.	N/A	N/A
10205771	12/2018	Palavalli et al.	N/A	N/A
10241820	12/2018	Lambeth et al.	N/A	N/A
10243797	12/2018	Lambeth et al.	N/A	N/A
10243834	12/2018	Shekhar et al.	N/A	N/A
10243846	12/2018	Jiang et al.	N/A	N/A
10243848	12/2018	Agarwal et al.	N/A	N/A
10257049	12/2018	Fried et al.	N/A	N/A
10298489	12/2018	Williams et al.	N/A	N/A
10333959	12/2018	Katrekar et al.	N/A	N/A
10339123	12/2018	Venkatesh et al.	N/A	N/A
10412018	12/2018	Feng et al.	N/A	N/A
10423790	12/2018	Patil et al.	N/A	N/A
10560343	12/2019	Cartsonis et al.	N/A	N/A
10579945	12/2019	Gaurav et al.	N/A	N/A
10587479	12/2019	Shakimov et al.	N/A	N/A
10601705	12/2019	Hira et al.	N/A	N/A
10616279	12/2019	Nimmagadda et al.	N/A	N/A
10637800	12/2019	Wang et al.	N/A	N/A
10673752	12/2019	Agarwal et al.	N/A	N/A
10693833	12/2019	Mathew et al.	N/A	N/A
10832224	12/2019	Palavalli et al.	N/A	N/A
10862753	12/2019	Hira et al.	N/A	N/A
10880158	12/2019	Lambeth et al.	N/A	N/A
10880170	12/2019	Wang et al.	N/A	N/A
10897420	12/2020	Pianigiani et al.	N/A	N/A
10908938	12/2020	Palavalli et al.	N/A	N/A
10942788	12/2020	Palavalli et al.	N/A	N/A
10999154	12/2020	Ahrenholz et al.	N/A	N/A

11057275	12/2020	Arunachalam et al.	N/A	N/A
11088902	12/2020	Palavalli et al.	N/A	N/A
11088916	12/2020	Chandrashekhhar et al.	N/A	N/A
11088919	12/2020	Chandrashekhhar et al.	N/A	N/A
11115301	12/2020	Margarian et al.	N/A	N/A
11133958	12/2020	Torvi et al.	N/A	N/A
11153170	12/2020	Chandrashekhhar et al.	N/A	N/A
11182163	12/2020	Beals et al.	N/A	N/A
11258668	12/2021	Chandrashekhhar et al.	N/A	N/A
11303557	12/2021	Chandrashekhhar et al.	N/A	N/A
11316773	12/2021	Dubey et al.	N/A	N/A
11336486	12/2021	Sharma et al.	N/A	N/A
11336556	12/2021	Chandrashekhhar et al.	N/A	N/A
11343227	12/2021	Vaidya et al.	N/A	N/A
11343283	12/2021	Vaidya et al.	N/A	N/A
11374817	12/2021	Chandrashekhhar et al.	N/A	N/A
11374850	12/2021	Chandrashekhhar et al.	N/A	N/A
11381456	12/2021	Manzanilla et al.	N/A	N/A
11394634	12/2021	Chandrashekhhar et al.	N/A	N/A
11438238	12/2021	Chandrashekhhar et al.	N/A	N/A
11496392	12/2021	Agarwal et al.	N/A	N/A
11509522	12/2021	Palavalli et al.	N/A	N/A
11528214	12/2021	Chandrashekar et al.	N/A	N/A
11601474	12/2022	Vaidya et al.	N/A	N/A
11683233	12/2022	Rogozinsky et al.	N/A	N/A
11736383	12/2022	Chandrashekhhar et al.	N/A	N/A
11743168	12/2022	Dubey et al.	N/A	N/A
11757940	12/2022	Vaidya et al.	N/A	N/A
2002/0029270	12/2001	Szczepanek	N/A	N/A
2002/0093952	12/2001	Gonda	N/A	N/A
2002/0131414	12/2001	Hadzic	N/A	N/A
2003/0046347	12/2002	Nishimura	N/A	N/A
2003/0167333	12/2002	Kumar et al.	N/A	N/A
2003/0185151	12/2002	Kurosawa et al.	N/A	N/A
2003/0185152	12/2002	Nederveen et al.	N/A	N/A
2003/0188114	12/2002	Lubbers et al.	N/A	N/A
2003/0188218	12/2002	Lubbers et al.	N/A	N/A
2004/0052257	12/2003	Abdo et al.	N/A	N/A
2005/0190757	12/2004	Sajassi	N/A	N/A

2005/0235352	12/2004	Staats et al.	N/A	N/A
2005/0288040	12/2004	Charpentier et al.	N/A	N/A
2006/0092976	12/2005	Lakshman et al.	N/A	N/A
2006/0179243	12/2005	Fields et al.	N/A	N/A
2006/0179245	12/2005	Fields et al.	N/A	N/A
2006/0193252	12/2005	Naseh et al.	N/A	N/A
2006/0195896	12/2005	Fulp et al.	N/A	N/A
2006/0221720	12/2005	Reuter	N/A	N/A
2006/0250249	12/2005	Cheng	N/A	N/A
2006/0251120	12/2005	Arimilli et al.	N/A	N/A
2007/0058631	12/2006	Mortier et al.	N/A	N/A
2007/0130295	12/2006	Rastogi et al.	N/A	N/A
2007/0217419	12/2006	Vasseur	N/A	N/A
2007/0219653	12/2006	Martin	N/A	N/A
2007/0239987	12/2006	Hoole et al.	N/A	N/A
2008/0013474	12/2007	Nagarajan et al.	N/A	N/A
2008/0049646	12/2007	Lu	N/A	N/A
2008/0104302	12/2007	Carpio	N/A	N/A
2008/0133729	12/2007	Fridman et al.	N/A	N/A
2008/0268847	12/2007	Mukherjee et al.	N/A	N/A
2008/0301379	12/2007	Pong	N/A	N/A
2009/0037367	12/2008	Wein	N/A	N/A
2009/0070337	12/2008	Romem et al.	N/A	N/A
2009/0193297	12/2008	Williams et al.	N/A	N/A
2009/0241192	12/2008	Thomas	N/A	N/A
2009/0279536	12/2008	Unbehagen et al.	N/A	N/A
2009/0279545	12/2008	Moonen	N/A	N/A
2009/0292858	12/2008	Lambeth et al.	N/A	N/A
2009/0296726	12/2008	Snively et al.	N/A	N/A
2010/0250784	12/2009	Henry et al.	N/A	N/A
2010/0257263	12/2009	Casado et al.	N/A	N/A
2010/0275199	12/2009	Smith et al.	N/A	N/A
2010/0322255	12/2009	Hao et al.	N/A	N/A
2011/0032898	12/2010	Kazmi et al.	N/A	N/A
2011/0047218	12/2010	Nojima et al.	N/A	N/A
2011/0051714	12/2010	Somes	N/A	N/A
2011/0085569	12/2010	Gnanasekaran et al.	N/A	N/A
2011/0164752	12/2010	Wainner et al.	N/A	N/A
2011/0188509	12/2010	Kern et al.	N/A	N/A
2011/0225123	12/2010	D'Souza	707/634	H04L 67/1034
2011/0231602	12/2010	Woods et al.	N/A	N/A
2011/0299413	12/2010	Chatwani et al.	N/A	N/A
2012/0084406	12/2011	Kumbalimutt	N/A	N/A
2012/0120964	12/2011	Koponen et al.	N/A	N/A
2012/0147898	12/2011	Koponen et al.	N/A	N/A
2012/0275328	12/2011	Iwata et al.	N/A	N/A
2013/0018947	12/2012	Archer et al.	N/A	N/A
2013/0024579	12/2012	Zhang et al.	N/A	N/A
2013/0042242	12/2012	Kagan	N/A	N/A

2013/0044636	12/2012	Koponen et al.	N/A	N/A
2013/0044641	12/2012	Koponen et al.	N/A	N/A
2013/0044761	12/2012	Koponen et al.	N/A	N/A
2013/0058250	12/2012	Casado et al.	N/A	N/A
2013/0058335	12/2012	Koponen et al.	N/A	N/A
2013/0058350	12/2012	Fulton	N/A	N/A
2013/0058354	12/2012	Casado et al.	N/A	N/A
2013/0058358	12/2012	Fulton et al.	N/A	N/A
2013/0060819	12/2012	Lambeth et al.	N/A	N/A
2013/0060940	12/2012	Koponen et al.	N/A	N/A
2013/0074065	12/2012	McNeeney et al.	N/A	N/A
2013/0103817	12/2012	Koponen et al.	N/A	N/A
2013/0125120	12/2012	Zhang et al.	N/A	N/A
2013/0132533	12/2012	Padmanabhan et al.	N/A	N/A
2013/0144992	12/2012	Barabash et al.	N/A	N/A
2013/0159637	12/2012	Forgette et al.	N/A	N/A
2013/0212212	12/2012	Addepalli et al.	N/A	N/A
2013/0215769	12/2012	Beheshti-Zavareh et al.	N/A	N/A
2013/0254328	12/2012	Inoue et al.	N/A	N/A
2013/0286833	12/2012	Torres et al.	N/A	N/A
2013/0287026	12/2012	Davie	N/A	N/A
2013/0301425	12/2012	Udutha et al.	N/A	N/A
2013/0301501	12/2012	Olvera-Hernandez et al.	N/A	N/A
2013/0301646	12/2012	Bogdanski et al.	N/A	N/A
2013/0308641	12/2012	Ackley	N/A	N/A
2014/0006354	12/2013	Parkison et al.	N/A	N/A
2014/0006357	12/2013	Davis et al.	N/A	N/A
2014/0006465	12/2013	Davis et al.	N/A	N/A
2014/0007239	12/2013	Sharpe et al.	N/A	N/A
2014/0064104	12/2013	Nataraja et al.	N/A	N/A
2014/0136908	12/2013	Maggiari et al.	N/A	N/A
2014/0146817	12/2013	Zhang	N/A	N/A
2014/0169365	12/2013	Sundaram et al.	N/A	N/A
2014/0172740	12/2013	McCormick et al.	N/A	N/A
2014/0201218	12/2013	Catalano et al.	N/A	N/A
2014/0208150	12/2013	Abuelsaad et al.	N/A	N/A
2014/0211661	12/2013	Gorkemli et al.	N/A	N/A
2014/0241356	12/2013	Zhang et al.	N/A	N/A
2014/0250220	12/2013	Kapadia et al.	N/A	N/A
2014/0269435	12/2013	McConnell et al.	N/A	N/A
2014/0301391	12/2013	Krishnan et al.	N/A	N/A
2014/0304355	12/2013	Kamath et al.	N/A	N/A
2014/0307744	12/2013	Dunbar et al.	N/A	N/A
2014/0337500	12/2013	Lee	N/A	N/A
2014/0351396	12/2013	Stabile et al.	N/A	N/A
2015/0009797	12/2014	Koponen et al.	N/A	N/A
2015/0010012	12/2014	Koponen et al.	N/A	N/A
2015/0016276	12/2014	Decusatis et al.	N/A	N/A

2015/0019444	12/2014	Purves	N/A	N/A
2015/0063364	12/2014	Thakkar et al.	N/A	N/A
2015/0067688	12/2014	Nagasawa et al.	N/A	N/A
2015/0085862	12/2014	Song	N/A	N/A
2015/0089034	12/2014	Stickler et al.	N/A	N/A
2015/0100704	12/2014	Davie et al.	N/A	N/A
2015/0103842	12/2014	Chandrashekhara et al.	N/A	N/A
2015/0103843	12/2014	Chandrashekhara et al.	N/A	N/A
2015/0106804	12/2014	Chandrashekhara et al.	N/A	N/A
2015/0117256	12/2014	Sabaa et al.	N/A	N/A
2015/0121483	12/2014	Perez et al.	N/A	N/A
2015/0180909	12/2014	Nakamatsu et al.	N/A	N/A
2015/0229641	12/2014	Sun et al.	N/A	N/A
2015/0234668	12/2014	Ravinoothala et al.	N/A	N/A
2015/0237014	12/2014	Bansal et al.	N/A	N/A
2015/0263946	12/2014	Tubaltsev et al.	N/A	N/A
2015/0264135	12/2014	Kandula et al.	N/A	N/A
2015/0312092	12/2014	Golshan	709/225	H04L 67/1023
2015/0312274	12/2014	Bishop et al.	N/A	N/A
2015/0312326	12/2014	Archer et al.	N/A	N/A
2015/0324220	12/2014	Bugenhagen	N/A	N/A
2015/0326467	12/2014	Fullbright et al.	N/A	N/A
2015/0381494	12/2014	Cherian et al.	N/A	N/A
2016/0057014	12/2015	Thakkar et al.	N/A	N/A
2016/0094396	12/2015	Chandrashekhara et al.	N/A	N/A
2016/0094661	12/2015	Jain et al.	N/A	N/A
2016/0134528	12/2015	Lin et al.	N/A	N/A
2016/0191374	12/2015	Singh et al.	N/A	N/A
2016/0210209	12/2015	Verkaik et al.	N/A	N/A
2016/0226700	12/2015	Zhang et al.	N/A	N/A
2016/0226762	12/2015	Zhang et al.	N/A	N/A
2016/0226822	12/2015	Zhang et al.	N/A	N/A
2016/0226959	12/2015	Zhang et al.	N/A	N/A
2016/0226967	12/2015	Zhang et al.	N/A	N/A
2016/0277289	12/2015	Madabushi et al.	N/A	N/A
2016/0352588	12/2015	Subbarayan et al.	N/A	N/A
2016/0373530	12/2015	Duda	N/A	N/A
2016/0380815	12/2015	Agarwal et al.	N/A	N/A
2016/0380891	12/2015	Agarwal et al.	N/A	N/A
2016/0380925	12/2015	Agarwal et al.	N/A	N/A
2016/0380973	12/2015	Sullenberger et al.	N/A	N/A
2017/0005988	12/2016	Bansal et al.	N/A	N/A
2017/0034052	12/2016	Chanda et al.	N/A	N/A
2017/0034129	12/2016	Sawant et al.	N/A	N/A
2017/0041347	12/2016	Nagaratnam et al.	N/A	N/A

2017/0048110	12/2016	Wu et al.	N/A	N/A
2017/0048130	12/2016	Goliya et al.	N/A	N/A
2017/0054758	12/2016	Maino	N/A	H04L 45/64
2017/0063633	12/2016	Goliya et al.	N/A	N/A
2017/0063794	12/2016	Jain et al.	N/A	N/A
2017/0063822	12/2016	Jain et al.	N/A	N/A
2017/0093617	12/2016	Chanda et al.	N/A	N/A
2017/0093636	12/2016	Chanda et al.	N/A	N/A
2017/0104720	12/2016	Bansal et al.	N/A	N/A
2017/0126431	12/2016	Han et al.	N/A	N/A
2017/0126497	12/2016	Dubey et al.	N/A	N/A
2017/0126551	12/2016	Pfaff et al.	N/A	N/A
2017/0163442	12/2016	Shen et al.	N/A	N/A
2017/0163532	12/2016	Tubaltsev et al.	N/A	N/A
2017/0163598	12/2016	Shen et al.	N/A	N/A
2017/0163599	12/2016	Shen et al.	N/A	N/A
2017/0171055	12/2016	Wang et al.	N/A	N/A
2017/0249195	12/2016	Sadana et al.	N/A	N/A
2017/0250912	12/2016	Chu et al.	N/A	N/A
2017/0264483	12/2016	Lambeth et al.	N/A	N/A
2017/0288981	12/2016	Hong et al.	N/A	N/A
2017/0289033	12/2016	Singh et al.	N/A	N/A
2017/0302531	12/2016	Maes	N/A	N/A
2017/0310641	12/2016	Jiang et al.	N/A	N/A
2017/0317954	12/2016	Masurekar et al.	N/A	N/A
2017/0317969	12/2016	Masurekar et al.	N/A	N/A
2017/0317971	12/2016	Dubey et al.	N/A	N/A
2017/0318113	12/2016	Ganichev et al.	N/A	N/A
2017/0324645	12/2016	Johnsen et al.	N/A	N/A
2017/0331711	12/2016	Duda	N/A	N/A
2017/0344444	12/2016	Costa-Roberts et al.	N/A	N/A
2018/0007004	12/2017	Basler	N/A	N/A
2018/0062880	12/2017	Yu et al.	N/A	N/A
2018/0062881	12/2017	Chandrashekhhar et al.	N/A	N/A
2018/0062923	12/2017	Katrekar	N/A	H04L 63/10
2018/0062944	12/2017	Altman et al.	N/A	N/A
2018/0063036	12/2017	Chandrashekhhar et al.	N/A	N/A
2018/0063195	12/2017	Nimmagadda et al.	N/A	N/A
2018/0083835	12/2017	Cole et al.	N/A	N/A
2018/0123877	12/2017	Saxena et al.	N/A	N/A
2018/0123903	12/2017	Holla et al.	N/A	N/A
2018/0157537	12/2017	Chen et al.	N/A	N/A
2018/0191682	12/2017	Liu et al.	N/A	N/A
2018/0234337	12/2017	Goliya et al.	N/A	N/A
2018/0234459	12/2017	Kung et al.	N/A	N/A
2018/0254947	12/2017	Baj	N/A	H04L 41/0894
2018/0309718	12/2017	Zuo	N/A	N/A

2018/0365095	12/2017	Monk	N/A	H04L 41/065
2018/0373961	12/2017	Wang et al.	N/A	N/A
2019/0014040	12/2018	Yerrapureddy et al.	N/A	N/A
2019/0069335	12/2018	Wu	N/A	N/A
2019/0109669	12/2018	Zachman et al.	N/A	N/A
2019/0158537	12/2018	Miriyala	N/A	N/A
2019/0190780	12/2018	Wang et al.	N/A	N/A
2019/0207847	12/2018	Agarwal et al.	N/A	N/A
2019/0245888	12/2018	Martinez et al.	N/A	N/A
2019/0260610	12/2018	Dubey et al.	N/A	N/A
2019/0260630	12/2018	Stabile et al.	N/A	N/A
2019/0297114	12/2018	Panchalingam	N/A	H04L 63/00
2019/0303326	12/2018	Desai et al.	N/A	N/A
2019/0334765	12/2018	Jain et al.	N/A	N/A
2019/0342175	12/2018	Wan	N/A	H04L 41/0816
2019/0363975	12/2018	Djernaes	N/A	N/A
2019/0379731	12/2018	Johnsen et al.	N/A	N/A
2020/0007392	12/2019	Goyal	N/A	N/A
2020/0007495	12/2019	Balamurugan	N/A	H04L 61/2567
2020/0007582	12/2019	Dixit et al.	N/A	N/A
2020/0007584	12/2019	Dixit et al.	N/A	N/A
2020/0014662	12/2019	Chanda et al.	N/A	N/A
2020/0021541	12/2019	Chanda	N/A	N/A
2020/0057669	12/2019	Hutcheson et al.	N/A	N/A
2020/0076684	12/2019	Naveen et al.	N/A	N/A
2020/0106744	12/2019	Miriyala et al.	N/A	N/A
2020/0162325	12/2019	Desai et al.	N/A	N/A
2020/0162337	12/2019	Jain et al.	N/A	N/A
2020/0169496	12/2019	Goliya et al.	N/A	N/A
2020/0169516	12/2019	Patel et al.	N/A	N/A
2020/0195607	12/2019	Wang et al.	N/A	N/A
2020/0257549	12/2019	Vlaznev et al.	N/A	N/A
2020/0296035	12/2019	Agarwal et al.	N/A	N/A
2020/0304427	12/2019	Sandler et al.	N/A	N/A
2020/0358693	12/2019	Rawlins	N/A	N/A
2020/0366741	12/2019	Kancherla et al.	N/A	N/A
2020/0409563	12/2019	Parasnis et al.	N/A	N/A
2021/0036889	12/2020	Jain et al.	N/A	N/A
2021/0067556	12/2020	Tahan	N/A	N/A
2021/0117908	12/2020	Coles et al.	N/A	N/A
2021/0126641	12/2020	Yan et al.	N/A	N/A
2021/0168197	12/2020	Jones et al.	N/A	N/A
2021/0194729	12/2020	Semwal et al.	N/A	N/A
2021/0311960	12/2020	Rogozinsky et al.	N/A	N/A
2021/0314192	12/2020	Chandrashekhhar et al.	N/A	N/A
2021/0314193	12/2020	Chandrashekhhar et al.	N/A	N/A

2021/0314212	12/2020	Chandrashekhhar et al.	N/A	N/A
2021/0314215	12/2020	Manzanilla et al.	N/A	N/A
2021/0314219	12/2020	Gujar et al.	N/A	N/A
2021/0314225	12/2020	Chandrashekhhar et al.	N/A	N/A
2021/0314226	12/2020	Chandrashekhhar et al.	N/A	N/A
2021/0314227	12/2020	Chandrashekhhar et al.	N/A	N/A
2021/0314228	12/2020	Chandrashekhhar et al.	N/A	N/A
2021/0314235	12/2020	Chandrashekhhar et al.	N/A	N/A
2021/0314251	12/2020	Dubey et al.	N/A	N/A
2021/0314256	12/2020	Chandrashekhhar et al.	N/A	N/A
2021/0314257	12/2020	Chandrashekhhar et al.	N/A	N/A
2021/0314258	12/2020	Chandrashekhhar et al.	N/A	N/A
2021/0314289	12/2020	Chandrashekhhar et al.	N/A	N/A
2021/0314291	12/2020	Chandrashekhhar et al.	N/A	N/A
2021/0367834	12/2020	Palavalli et al.	N/A	N/A
2022/0103429	12/2021	Vaidya et al.	N/A	N/A
2022/0103430	12/2021	Vaidya et al.	N/A	N/A
2022/0103514	12/2021	Vaidya et al.	N/A	N/A
2022/0103521	12/2021	Vaidya et al.	N/A	N/A
2022/0103598	12/2021	Vaidya et al.	N/A	N/A
2022/0191126	12/2021	Dubey et al.	N/A	N/A
2022/0255896	12/2021	Jain et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
3119423	12/2017	CA	N/A
102124456	12/2010	CN	N/A
102986172	12/2012	CN	N/A
103650433	12/2013	CN	N/A
103890751	12/2013	CN	N/A
110061899	12/2018	CN	N/A
1154601	12/2000	EP	N/A
1290856	12/2002	EP	N/A
1635506	12/2005	EP	N/A
1868318	12/2006	EP	N/A
3016331	12/2015	EP	N/A
3485610	12/2018	EP	N/A
4078395	12/2021	EP	N/A

2010028364	12/2009	WO	N/A
2011140028	12/2010	WO	N/A
2012113444	12/2011	WO	N/A
2013026049	12/2012	WO	N/A
2013152716	12/2012	WO	N/A
2015054671	12/2014	WO	N/A
2017003881	12/2016	WO	N/A
2018044341	12/2017	WO	N/A
2021206785	12/2020	WO	N/A
2021206786	12/2020	WO	N/A
2021206790	12/2020	WO	N/A
2022066269	12/2021	WO	N/A

OTHER PUBLICATIONS

Author Unknown, "What is a Data Center?," Cyberpedia, Month Unknown 2022, 5 pages, Palo Alto Networks, retrieved from <https://www.paloaltonetworks.com/cyberpedia/what-is-a-data-center>. cited by applicant

Author Unknown, "Apache Cassandra™ 1.2 Documentation," Jan. 13, 2013, 201 pages, DataStax. cited by applicant

Author Unknown, "OpenFlow Switch Specification, Version 1.1.0 Implemented (Wire Protocol 0x02)," Feb. 28, 2011, 56 pages, Open Networking Foundation. cited by applicant

Berde, Pankaj, et al., "ONOS Open Network Operating System An Open-Source Distributed SDN OS," Dec. 19, 2013, 34 pages. cited by applicant

Giannakou, Anna, et al., "AL-SAFE: A Secure Self-Adaptable Application-Level Firewall for IaaS Clouds," 2016 IEEE 8th International Conference on Cloud Computing Technology and Science, Dec. 12-15, 2016, 8 pages, IEEE, Luxembourg City, LU. cited by applicant

Hanna, Jeremy, "How ZooKeeper Handles Failure Scenarios," <http://apache.org/hadoop/Zookeeper/FailureScenarios>. Dec. 9, 2010, 1 page. cited by applicant

Heller, Brandon, et al., "The Controller Placement Problem," Hot Topics in Software Defined Networks, Aug. 13, 2012, 6 pages, Helsinki, Finland. cited by applicant

Kohila, N., et al., "Data Security in Local Network Using Distributed Firewall," International Journal of Scientific Research in Computer Science Applications and Management Studies, Nov. 2014, 8 pages, vol. 3, Issue 6, jsrscams.com. cited by applicant

Krishnaswamy, Umesh, et al., "ONOS Open Network Operating System—An Experimental Open-Source Distributed SDN OS," Apr. 16, 2013, 24 pages. cited by applicant

Lebresne, Sylvain, "[Release] Apache Cassandra 1.2 released," Jan. 2, 2013, 1 page. cited by applicant

Loshin, Peter, et al., "What is a data center?," Special Report: Everything You Need to Know About the Log4j Vulnerability, Oct. 2021, 13 pages, retrieved from <https://searchdatacenter.techtarget.com/definition/data-center>. cited by applicant

Mahalingham, Mallik, et al., "VXLAN: A Framework for Overlaying Virtualized Layer 2 Networks over Layer 3 Networks," draft-mahalingham-dutt-dcops-vxlan-02.txt Internet Draft, Aug. 22, 2012, 20 pages, Internet Engineering Task Force. cited by applicant

Mahalingham, Mallik, et al., "VXLAN: A Framework for Overlaying Virtualized Layer 2 Networks over Layer 3 Networks," draft-mahalingham-dutt-dcops-vxlan-00.txt Internet Draft, Aug. 26, 2011, 20 pages, Internet Engineering Task Force. cited by applicant

Non-Published Commonly Owned U.S. Appl. No. 18/227,655, filed Jul. 28, 2023, 133 pages, VMware, Inc. cited by applicant

PCT International Search Report and Written Opinion of Commonly Owned International Patent Application PCT/US2021/015968, mailed Apr. 23, 2021, 14 pages, International Searching

Authority (EP). cited by applicant
Schuba, Christoph, et al., "Integrated Network Service Processing Using Programmable Network Devices," SMLI TR-2005-138, May 2005, 30 pages, Sun Microsystems, Inc. cited by applicant
Surantha, Nico, et al., "Secure Kubernetes Networking Design Based on Zero Trust Model: A Case Study of Financial Service Enterprise in Indonesia," Innovative Mobile and Internet Services in Ubiquitous Computing, Jan. 2020, 14 pages, Springer Nature, Switzerland. cited by applicant
Wack, John, et al., "Guidelines on Firewalls and Firewall Policy," Recommendations of the National Institute of Standards and Technology, Special Publication 800-41, Jan. 2002, 75 pages, NIST, Gaithersburg, MD, USA. cited by applicant

Primary Examiner: Raza; Muhammad

Attorney, Agent or Firm: Quarles & Brady LLP

Background/Summary

CLAIM OF BENEFIT TO PRIOR APPLICATIONS (1) This application is a continuation application of U.S. patent application Ser. No. 16/906,955, filed Jun. 19, 2020, now published as U.S. Patent Publication 2021/0314227. U.S. patent application Ser. No. 16/906,955 claims the benefit of Indian Patent Application number 202041015134, filed Apr. 6, 2020. U.S. patent application Ser. No. 16/906,955, now published as U.S. Patent Publication 2021/0314227, is incorporated herein by reference.

BACKGROUND

(1) As more networks move to the cloud, it is more common for corporations or other entities to have networks spanning multiple sites. While logical networks that operate within a single site are well established, there are various challenges in having logical networks span multiple physical sites (e.g., datacenters). The sites should be self-contained, while also allowing for data to be sent from one site to another easily. Various solutions are required to solve these issues.

BRIEF SUMMARY

(2) Some embodiments of the invention provide a network management system for managing a logical network spanning multiple federated sites (e.g., multiple datacenters). The network management system of some embodiments includes a global manager that manages the entire logical network spanning all of the sites, as well as local managers at each site that directly manage the logical network at their respective sites. The logical network includes logical network elements that span one or more sites and logical network policies that apply to the elements at those sites. In some embodiments, the global manager receives a global desired configuration for the logical network (e.g., from an administrator of the network), identifies a relevant portion of the global desired configuration for each site in the federation, and provides the identified portion to the site's corresponding local manager. In addition, the global manager collects information about the realized state of the logical network across the sites from the local managers and provides this information (e.g., to an administrator of the logical network) for troubleshooting and management purposes.

(3) In some embodiments, the global manager executes on a computing device at one of the sites spanned by the logical network, and each local manager also executes on a computing device at its respective site. In some embodiments, the global manager executes on the same computing device at one of the physical sites as the local manager managing that site. The global manager and the local managers are in some embodiments separate modules of a single application. Some embodiments deploy each manager at a physical site as a cluster of machines, with each machine

executing on a different computing device at the same site.

(4) Some embodiments employ a primary global manager and a secondary global manager, in an active-standby arrangement. The primary global manager is asynchronously synchronized with the secondary global manager as a standby for failover scenarios. The secondary global manager executes on a different computing device, located in some embodiments for resiliency at a different site than the primary global manager, and maintains an independent database. The secondary global manager is also deployed in some embodiments as a set of machines and may also execute on the same computing device as a local manager.

(5) The primary global manager's database is a distributed shared log implemented across a set of storage devices at the physical site where the primary global manager resides in some embodiments. Data regarding the global desired configuration is received and stored in the database using a series of database transactions which are initiated through a series of application programming interface (API) calls to the global manager. The database, in some embodiments, generates an update stream from the database transactions, that is used to provide the data regarding the desired configuration to the secondary global manager for replication of the database. The update stream also includes metadata associated with each transaction, such as timestamp information that can be used for data ordering, as well as database status to prevent race conditions for access. In some embodiments, the database is shared by the primary global manager with other applications (e.g., a local manager) on the same computing device. In some such embodiments, data for replication to the secondary global manager is tagged so that only data associated with the primary global manager is replicated and other data associated with other applications on the computing device is not replicated.

(6) Each global manager's database also includes in some embodiments a set of queues, each of which is reserved for a local manager at each of the physical sites, as well as one queue reserved for the secondary global manager. The database at the secondary global manager also includes these queues. When the primary global manager receives the global desired configuration for the logical network, the global manager stores portions of the global configuration in each queue, based on the relevance of the portions to the configuration of the logical network at the queue's corresponding physical site. In some embodiments, a broker service of the global manager identifies the relevant portions for each physical site (e.g., based on the span of the logical network elements), as described in further detail below.

(7) In some embodiments, a set of asynchronous channels connect the primary global manager to the local managers and the secondary global manager at different sites. Some embodiments execute a set of site managers at each global manager to manage the channels, with each site manager corresponding to one of the physical sites. The channels retrieve data from the queues in the database and provide the retrieved data to the destination database at the corresponding physical site. These channels maintain the connections between physical sites and in some embodiments guarantee various connection parameters (e.g., the minimum bandwidth, the maximum roundtrip time, etc.) that are required for replication of data to the secondary global manager and dissemination of data to the local managers. The channels also identify the active machine for each manager, when the managers are implemented by a cluster of machines.

(8) In some embodiments, the global desired configuration of the logical network is expressed as a hierarchical tree (also referred to as a global policy tree) with nodes and connections between the nodes. In some embodiments, the nodes represent logical network elements that span one or more sites and logical network policies that apply to those elements, and the connections represent relationships between the nodes (e.g., parent-child relationships, logical network connections, etc.). The logical network elements include logical forwarding elements that forward data in the logical network, e.g. logical routers, logical switches, etc. These logical network elements also include logical constructs in some embodiments, such as logical ports associated with the logical forwarding elements, domains that are logical groupings of one or more sites (e.g., geographic

regions), and groups of logical network endpoints that share one or more attributes (e.g., operating system, region, etc.). Each logical network element is implemented in some embodiments by physical forwarding elements (PFEs) executing on computing devices at the sites that are spanned by that logical network element. The logical network policies include forwarding policies, service policies, and security policies, and are applied in some embodiments to govern the behavior of the logical forwarding elements.

(9) The global manager performs a span calculation in some embodiments for each logical network element, to determine which logical network elements are implemented at each site. In some embodiments, the span calculation is based on the relationships between the nodes in the global policy tree. For example, when the relationship between two nodes is a parent-child relationship, the child node may inherit its span from the parent node. In other cases, however, a child node does not inherit the entire span of its parent node. As another example, when the relationship between two nodes is a dependence of one node on another node, expressed as a reference from one node to the other, the span of the dependent node will depend on the referenced node. Some logical network elements also have a pre-defined span in some embodiments, e.g., defined by an administrator of the network.

(10) The relevant portion of the global desired configuration that is provided to each local manager in some embodiments is a site-specific policy subtree of the global policy tree, corresponding to the logical network elements implemented at that site. The subtree for each site only includes nodes for logical network elements that span the site. In some embodiments, the subtree is generated by first identifying the span of each node in the global policy tree. The global policy tree is then parsed, using the identified span for each node, to generate the policy subtree for each site.

(11) The local manager at each site uses the relevant portion of the global desired configuration, received from the global manager, to manage the logical network at the site. For example, in some embodiments, the local manager uses the relevant portion to generate and provide configuration data to the control plane of the logical network (e.g., a cluster of controllers at each site). In some embodiments, these controllers identify computing devices at the site which execute physical forwarding elements, and distribute the configuration data to the identified computing devices. Some embodiments have local controllers (also referred to as chassis controllers) that execute on one or more of the computing devices, and which receive the configuration data from the controller cluster. The local controllers use the configuration data to configure the physical forwarding elements to implement the logical network elements. Each site's controller cluster also creates mappings between logical addresses (e.g., MAC addresses of logical network endpoints executing on the computing devices) and physical addresses (e.g., IP addresses of tunnel endpoints at the computing devices), and distributes these mappings to each computing device to which they are relevant, as well as to other controller clusters at other sites that require the data.

(12) In some embodiments, the computing devices at each site also execute machines along with the physical forwarding elements and the local controllers. These machines include logical network endpoints, which are sources and destinations of data message traffic (e.g. computers, virtual machines, containers, etc.), and service machines, which perform services on the data traffic (e.g. firewalls, load balancers, etc.). A machine is located at a single site at a given time, but may be migrated between sites in some embodiments. These machines connect in some embodiments to the physical forwarding elements in order to exchange data messages in the network. In some embodiments, logical network endpoints at each site are logically organized into groups which can span multiple sites. The service machines in some embodiments apply some of the logical network policies to the data messages exchanged between groups of endpoints based on policies that are defined in terms of these groups.

(13) Policies are defined in some embodiments at the global manager, e.g., by an administrator of the logical network. In some embodiments, policies are one or more service rules which are enforced at the sites on data message flows based on a set of flow attributes. The global manager in

some embodiments distributes the service rules to local managers at the sites at which the policy is to be enforced (i.e., the policy's span). The local manager uses the service rules to generate configuration data for distribution by controllers, to configure the data plane (i.e., the forwarding elements and the service machines) to enforce the received service rules on data message flows that are associated with groups of logical network endpoints.

(14) The policies are defined in some embodiments by reference to groups of logical network endpoints that span one or more sites. The groups are defined in some embodiments at the global manager, e.g. by an administrator of the logical network. The service rules refer to these groups in some embodiments by using a group identifier that is assigned at the global manager when the groups are defined. The definitions of these groups are distributed to the sites spanned by the policy, so that the controllers at these sites are able to configure the data plane to enforce the policy.

(15) As discussed above, the service rules for a given policy are enforced on data messages that are associated with a group of logical network endpoints spanning one or more sites. Service rules have multiple fields in some embodiments, including source fields, destination fields, and action fields. Some embodiments refer to these groups in these fields by using group identifiers that are assigned to the groups (e.g., by the global manager when the groups are defined). For example, when the group identifier for the group is referenced in a source field of the service rule, the service rule is enforced on data messages that originate from machines in that group. When the group identifier for the group is referenced in a destination field of the service rule, the service rule is enforced on data messages that are directed to machines in that group.

(16) The service rule is enforced on data messages whose attributes match those specified in the rule. For example, the rule may specify a flow 5-tuple (source IP address and port number, destination IP address and port number, and transmission protocol). Alternatively, the rule may specify other attributes of the flow. The PFEs at each site identify matching flows by performing match classification operations on each data message, to determine if the flow attributes match those specified in the service rule. When a data message matches the service rule, it is provided to a service machine in some embodiments, which performs the action that is specified by the rule on the matching data message. These actions include dropping the data message, allowing the data message, or performing other service operations on the data message. For example, a service operation may modify the header of the data message, to redirect the message to a different destination. These service operations include load-balancing, firewall, Dynamic Host Configuration Protocol (DHCP), Network Address Translation (NAT), and other services.

(17) Some embodiments define domains that are groups of one or more sites. The domains are defined in some embodiments at the global manager, e.g. by an administrator of the logical network. Some domains are specific to a single physical site, and are referred to as locations. Other domains are logical groups of one or more sites, and are referred to as regions. Some embodiments restrict each site to membership in a single location and a single region. In other words, a location may not have more than one site, and a site may not be in more than two regions. The member sites of each domain define the domain's span. In some embodiments, domains are defined at the global manager and represented as nodes in the policy tree. As noted above, policies are also defined as nodes in the policy tree, and in some embodiments, policies are defined as child nodes of domain nodes. In such cases, the policy (and therefore, any service rule of the policy) inherits the span of its parent domain, i.e. the policy is enforced at the sites which are members of its parent domain.

(18) Furthermore, groups are also represented in some embodiments as child nodes of a domain. In such cases, the group inherits the span of its parent domain, i.e. the group is defined at the sites which are members of its parent domain. Logical network endpoints located at a site may join a group if the group is defined at that site (i.e., if the group spans that site).

(19) In some embodiments, a policy's service rules distributed to a first set of sites for applying to a first group of machines defined at those sites refer to a second group of machines defined at a second, different set of sites. In other words, the span of the referenced group does not include the

span of the policy. Since such service rules reference a group of machines that is undefined for one or more of the sites in the first set of sites, not all of the machines in the first group are able to enforce the policy. Therefore, some embodiments identify which sites in the first set of sites do not have the definition of the referenced group, and distribute the definition of the referenced group to those identified sites.

(20) For example, the policy might be defined in a first domain, which has the first set of sites as members, while the group referenced by the policy is defined in a second, different domain, which has the second set of sites as members. The policy therefore inherits span from the first domain, and the referenced group inherits span from the second domain. In such embodiments, the definition of the referenced group is distributed to the identified sites by extending the span of the group to include the span of the policy. This is done at the global manager (e.g., by an administrator of the logical network) in some embodiments, by creating a reference group in the first domain which references the group in the second domain. Upon receiving this modification, the global manager repeats the span calculation and determines that the group now spans the first set of sites. The global manager then provides the portion of the desired configuration (which now includes the referenced group) to the local managers at the first set of sites, which are now able to configure their respective controllers to configure the physical forwarding elements and service machines at the first set of sites to enforce the rule.

(21) Some embodiments provide, upon request, the realized state of one or more of the logical network elements of the logical network. In some embodiments, the global manager receives the request for the realized state of the logical elements from an administrator of the logical network. For each logical network element in the request, the global manager identifies the sites spanned by the element and sends requests to those identified sites' local managers. These requests are sent in parallel in some embodiments. After receiving the realized state from each of the sites spanned by each of the logical network elements, the global manager combines the received realized state from each site and presents the realized state in response to the initial request. For example, the realized state may be presented in a report for an administrator of the logical network in some embodiments.

(22) The realized state of a logical network element indicates in some embodiments whether the desired configuration (also referred to as the desired state) of the logical network element has been realized at each of the sites that the element spans. Some embodiments describe the realized state partly in terms of a deviation from the desired state. In some embodiments, the realized state for a logical network element (e.g., a logical router, logical switch, etc.) describes how the logical network element is implemented at each of the physical sites that it spans. For example, the realized state includes data from PFEs executing on computing devices at each of the spanned sites to implement the logical network element. Identifying the spanned sites is achieved in some embodiments by performing a span calculation as described above, or determined from the definition of the logical network element (e.g., by an administrator of the logical network). For each logical network element, the realized state received from each local manager is a subset of the total realized state of that logical network element, just as the desired state for the logical network element that is provided to the local manager is a subset of the global desired state of that logical network element defined at the global manager.

(23) In some embodiments, the global manager also receives notifications from the local managers. Unlike requests for realized state which are provided on-demand (e.g., pulled from the local managers), these notifications are automatically pushed from the local managers. The notifications pertain in some embodiments to site-specific implementation problems or situations. Notifications are provided in some embodiments to the GM as the events they describe occur (e.g., in real-time or near-real-time rather than on-demand). These notifications are also displayed in some embodiments as a report (e.g., for review by an administrator of the logical network).

(24) One type of notification in some embodiments is an alarm from a local manager that the

implementation of a logical network element has failed at the specific site. For example, the failure could be due to an IP address conflict with another network element defined locally at the local manager, of which the global manager was unaware.

(25) Another type of notification in some embodiments is a message from the local manager that an administrator of the network has overridden the desired state of a logical network element at that particular site. This occurs in some embodiments by an administrator (e.g., a different administrator, possibly located in a different region or even country) using the local manager to modify the logical network rather than using the global manager. Some embodiments restrict overriding the desired state to networking-related configurations only, and prevent such overrides for security-related configurations (e.g., security policies).

(26) A third type of notification in some embodiments is a message from the local manager that the deletion of a logical network element has failed at the site. In other words, the logical network element was deleted at the global manager (e.g. for being unneeded or unused), but at the local manager of one site, the logical network element is still in use. An example may be a first logical router defined locally at one site's local manager, and which connects to a second logical router defined at the global manager. The global manager is unaware of the first logical router, which does not exist at other sites, and a global administrator may attempt to delete the second logical router. The local manager at the site where the second logical router is in use as a link for the first logical router then sends a notification to inform the global manager that the second logical router has not yet been deleted from that site.

(27) An additional example of a type of notification in some embodiments is a message from the local manager that a logical network endpoint (e.g., a virtual machine) has been migrated from one site to another site. Such endpoints are typically attached to a logical port of a logical network element (e.g., a logical switch) that spans both sites. When the endpoint is migrated to a different site, it retains its association with that same logical port in some embodiments, and seamlessly inherits the same security policies that are associated with that port, even at the new site.

(28) In some embodiments, a machine may be migrated from one physical site to another physical site, while preserving the state of network policies that apply to the machine. For example, for a logical segment (logical switch) that spans at least two physical sites, a machine attached to a logical port of the segment at one of the spanned sites is migrated in some embodiments to a new logical port of the same segment at another one of the spanned sites. Migrating a machine in some embodiments comprises creating a new machine at a new host computer (at the new site), transferring configuration data to the new machine from the original machine, and removing the original machine from its host computer. In some embodiments, the configuration data comprises physical state data and logical network policies.

(29) The state of the machine (e.g., runtime data) is transferred in some embodiments from a computing device executing the original machine at the source site (i.e., the source host) to a computing device executing the new machine at the destination site (i.e., the destination host). In some embodiments, the state includes a local port identifier assigned by a local manager at the source physical site, and a global port identifier assigned by the global manager. The global port identifier is associated with logical network policies defined through a user client (e.g., by an administrator of the network) in some embodiments as part of the desired configuration of the logical network. Upon receiving the transferred state data, the destination host overwrites at least a portion of the state of the new machine with the received state of the old machine.

(30) In some embodiments, the destination host extracts the global port identifier from the received state data and provides it to the destination site's local manager. The local manager uses the global port identifier in some embodiments to retrieve logical network policies associated with the global port identifier from the global manager, and applies these policies to the new machine. In some embodiments, the local manager generates configuration data from the policies and provides the configuration data to the control plane of the logical network at the destination site (e.g., a set of

controllers) which then distribute the configuration data to the destination host.

(31) The preceding Summary is intended to serve as a brief introduction to some embodiments of the invention. It is not meant to be an introduction or overview of all of the inventive subject matter disclosed in this document. The Detailed Description that follows and the Drawings that are referred to in the Detailed Description will further describe the embodiments described in the Summary as well as other embodiments. Accordingly, to understand all the embodiments described by this document, a full review of the Summary, Detailed Description and the Drawings is needed. Moreover, the claimed subject matters are not to be limited by the illustrative details in the Summary, Detailed Description and the Drawings, but rather are to be defined by the appended claims, because the claimed subject matters can be embodied in other specific forms without departing from the spirit of the subject matters.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The novel features of the invention are set forth in the appended claims. However, for purposes of explanation, several embodiments of the invention are set forth in the following figures.
- (2) FIG. 1 conceptually illustrates a simple example of a logical network of some embodiments.
- (3) FIG. 2 conceptually illustrates the logical routing components of the logical routers.
- (4) FIG. 3 conceptually illustrates physical sites spanned by the logical network, with the host computers and edge devices that implement the logical network.
- (5) FIG. 4 conceptually illustrates an example of a network management system of some embodiments for a logical network that spans multiple physical sites.
- (6) FIG. 5 conceptually illustrates a global manager executing on a separate computing device at a given site from any computing device executing the site's local manager.
- (7) FIG. 6 conceptually illustrates a global manager executing on a computing device at a separate site that is not spanned by the logical network, and therefore has no local managers.
- (8) FIG. 7 conceptually illustrates different components of the network management system described in FIG. 4.
- (9) FIG. 8 conceptually illustrates a manager cluster and a distributed database system at a physical site spanned by the logical network.
- (10) FIG. 9 conceptually illustrates generating an update stream for use by the primary global manager, to replicate the desired configuration to the secondary global manager.
- (11) FIG. 10 illustrates a process performed in some embodiments by a database instance to generate an update stream.
- (12) FIG. 11 conceptually illustrates a broker service retrieving a create, update, or delete event from the database and populating relevant persistent queues.
- (13) FIG. 12 illustrates a process performed in some embodiments by the broker service to populate the persistent queues.
- (14) FIG. 13 conceptually illustrates a process performed in some embodiments by the sending AR module at the primary global manager for providing the retrieved data to the secondary global manager and the local managers.
- (15) FIG. 14 conceptually illustrates the architecture of a local manager receiving data through a channel from the primary global manager.
- (16) FIG. 15 illustrates a process performed in some embodiments by a local manager when receiving data from the global manager.
- (17) FIG. 16 conceptually illustrates a process performed in some embodiments by the local manager when it receives a CUD event directly from a user client, instead of from the global manager.

- (18) FIG. 17 conceptually illustrates an example of a global policy tree of some embodiments.
- (19) FIG. 18 conceptually illustrates examples of security groups that span one or more physical sites.
- (20) FIG. 19 conceptually illustrates a process performed in some embodiments by the global manager to parse the global desired configuration in order to populate the global policy tree and identify, for each physical site, a relevant portion of the global desired configuration.
- (21) FIG. 20 conceptually illustrates a process performed in some embodiments by the global manager to generate policy subtrees.
- (22) FIG. 21 conceptually illustrates an example of a global policy subtree for a single physical site.
- (23) FIG. 22 conceptually illustrates a local policy tree at a single physical site.
- (24) FIG. 23 conceptually illustrates a match-action table of some embodiments.
- (25) FIG. 24 conceptually illustrates a process performed in some embodiments by the global manager to extend the span of a group for applying a cross-domain policy.
- (26) FIGS. 25A-B conceptually illustrate an example of extending the group span due to a cross-domain policy.
- (27) FIG. 26 conceptually illustrates a process performed in some embodiments by the global manager to retrieve the realized state of a logical network element.
- (28) FIG. 27 conceptually illustrates a scatter-and-gather operation of the global manager to retrieve, on demand, the realized state of a logical network element that spans multiple physical sites.
- (29) FIGS. 28A-C conceptually illustrate an example of a user interface of some embodiments for presenting the realized state of a logical router to an administrator of the logical network.
- (30) FIGS. 29A-B conceptually illustrate an example of a user interface of some embodiments for presenting the realized state of a logical switch to an administrator of the logical network.
- (31) FIG. 30 conceptually illustrates an example of an alarm caused by a failure to create a requested logical network element at a local manager.
- (32) FIG. 31 conceptually illustrates an example of an alarm caused by an attempt to modify the global configuration.
- (33) FIGS. 32A-B conceptually illustrate an example of an alarm caused by a failure to delete a logical network element.
- (34) FIG. 33 illustrates an example of some embodiments for migrating a virtual machine between sites within a region while preserving the machine's group membership and associated policies.
- (35) FIGS. 34A-C conceptually illustrate the exchange of information during creation and migration of a virtual machine between the global manager, the local managers, and the source and destination hosts.
- (36) FIG. 35 conceptually illustrates a process performed in some embodiments by a source host during the creation and migration of a virtual machine.
- (37) FIG. 36 conceptually illustrates a process performed in some embodiments by a destination host during the creation and migration of a virtual machine.
- (38) FIG. 37 conceptually illustrates an electronic system with which some embodiments of the invention are implemented.

DETAILED DESCRIPTION

- (39) In the following detailed description of the invention, numerous details, examples, and embodiments of the invention are set forth and described. However, it will be clear and apparent to one skilled in the art that the invention is not limited to the embodiments set forth and that the invention may be practiced without some of the specific details and examples discussed.
- (40) Some embodiments of the invention provide a network management system for managing a federated logical network spanning multiple physical sites (e.g., datacenters). The desired configuration of the logical network is defined by one or more network administrators using a set

of network managers. The logical network includes logical network elements that span one or more sites and logical network policies that apply to the elements at those sites. Any such elements or policies that span multiple physical sites are defined through the global manager.

(41) In some embodiments, the global manager receives (e.g., via a user client) a global desired configuration for the logical network, identifies a relevant portion of the global desired configuration for each physical site, and provides the identified portion to the site's corresponding local manager. In addition, the global manager collects information about the realized state of the logical network across the sites from the local managers and provides this information (e.g., to an administrator of the logical network) for troubleshooting and management purposes.

(42) The logical network elements of some embodiments include logical forwarding elements (LFEs), such as logical switches (to which logical network endpoints attach) and logical routers. Each LFE (e.g., logical switch or logical router) is implemented across one or more physical sites, depending on how the LFE is defined by the network administrator. In some embodiments, the LFEs are implemented within the physical sites by physical forwarding elements (PFEs) executing on host computers that also host logical network endpoints of the logical network (e.g., with the PFEs executing in virtualization software of the host computers) and/or on edge devices within the physical sites. The LFEs transmit data traffic between logical network endpoints (e.g., data compute nodes (DCNs)) (i) within a datacenter, (ii) between logical network endpoints in different datacenters, and (iii) between logical network endpoints in a datacenter and endpoints external to the logical network (e.g., external to the datacenters). The edge devices, in some embodiments, are computing devices that may be bare metal machines executing a datapath and/or computers on which logical network endpoints execute to a datapath. These datapaths, in some embodiments, perform various gateway operations (e.g., gateways for stretching logical switches across physical sites, gateways for executing centralized features of logical routers such as performing stateful services and/or connecting to external networks).

(43) FIG. 1 conceptually illustrates a simple example of a logical network **100** of some embodiments. This logical network **100** includes a tier-0 (T0) logical router **105**, a tier-1 (T1) logical router **110**, and logical switches **115** and **120**. Though not shown, various logical network endpoints (e.g., virtual machines, containers, or other types of data compute nodes) attach to logical ports of the logical switches **115** and **120**. These logical network endpoints execute on host computers in the physical sites spanned by the logical switches to which they attach. In this example, both the T0 logical router **105** and the T1 logical router **110** are defined to span three physical sites, though such routers can span any number of sites in other embodiments. In some embodiments, the logical switches **115** and **120** inherit the span of the logical router **105** to which they connect.

(44) As in this example, logical routers, in some embodiments, may include T0 logical routers (e.g., router **105**) that connect directly to external networks **125** and T1 logical routers (e.g., router **110**) that segregate a set of logical switches from the rest of the logical network and may perform stateful services for endpoints connected to those logical switches. These logical routers **105-110**, in some embodiments, are defined by the network managers to have one or more routing components, depending on how the logical router has been configured by the network administrator.

(45) FIG. 2 conceptually illustrates the logical network **200** showing the logical routing components of the logical routers **105** and **110** as well as the various logical switches that connect to these logical components and that connect the logical components to each other. The logical routers **105-110** span three physical sites **201-203**, with some components of each router spanning all three sites and others only spanning a single site. As shown, the T1 logical router **110** includes a distributed routing component (DR) **205** as well as a set of centralized routing components (also referred to as service routers, or SRs) **210-220**. T1 logical routers, in some embodiments, may have only a DR, or may have both a DR as well as SRs. For T1 logical routers, SRs allow for centralized

(e.g., stateful) services to be performed on data messages sent between (i) logical network endpoints connected to logical switches that connect to the T1 logical router and (ii) logical network endpoints connected to other logical switches that do not connect to the tier-1 logical router or from external network endpoints. In this example, data messages sent to or from logical network endpoints connected to logical switches **115** and **120** will have stateful services applied by one of the SRs **210-220** of the T1 logical router **110** (specifically, by the primary SR **215**).

(46) T1 logical routers may be connected to T0 logical routers in some embodiments (e.g., T1 logical router **110** connecting to T0 logical router **105**). These T0 logical routers **105**, as mentioned, handle data messages exchanged between the logical network endpoints and external network endpoints. As shown, the T0 logical router **105** includes a DR **225** as well as a set of SRs **230-240**. In some embodiments, T0 logical routers include an SR (or multiple SRs) operating in each physical site spanned by the logical router. In some or all of these physical sites, the T0 SRs **230-240** connect to external routers **241-243** (or to top of rack (TOR) switches that provide connections to external networks).

(47) In addition to the logical switches **115** and **120** (which span all of the physical sites spanned by the T1 DR **205**), FIG. 2 also illustrates various automatically-defined logical switches. Within each physical site, the T1 DR **205** connects to its respective local T1 SR **210-220** via a respective transit logical switch **245-255**. Similarly, within each physical site, the T0 DR **225** connects to its respective local T0 SR **230-240** via a respective transit logical switch **260-270**. In addition, a router link logical switch **275** connects the primary T1 SR **215** (that performs the stateful services for the T1 logical router **110**) to the T0 DR **225**. In some embodiments, similar router link logical switches are defined for each of the other physical sites, but are marked as down.

(48) Lastly, the network management system also defines backplane logical switches that connect each set of SRs. In this case, there is a backplane logical switch **280** connecting the T1 SRs **210-220** and a backplane logical switch **285** connecting the T0 SRs **230-240**. These backplane logical switches, unlike the transit logical switches, are stretched across the physical sites spanned by their respective logical routers. When one SR for a particular logical router routes a data message to another SR for the same logical router, the data message is sent according to the appropriate backplane logical switch.

(49) As mentioned, the LFEs of a logical network may be implemented by PFEs executing on source host computers as well as by the edge devices. FIG. 3 conceptually illustrates physical sites **305-315** spanned by the logical network **300** with the host computers **320** and edge devices **325** that implement the logical network. Virtual machines (VMs) (in this example) or other logical network endpoints operate on the host computers **320**, which execute virtualization software for hosting these VMs. The virtualization software, in some embodiments, includes the PFEs such as virtual switches and/or virtual routers. In some embodiments, one PFE (e.g., a flow-based PFE) executes on each host computer **320** to implement multiple LFEs, while in other embodiments multiple PFEs execute on each host computer **320** (e.g., one or more virtual switches and/or virtual routers). In still other embodiments, different host computers execute different virtualization software with different types of PFEs. Within this application, “PFE” is used to represent the set of one or more PFEs that execute on a host computer to implement LFEs of one or more logical networks.

(50) The edge devices **325**, in some embodiments, execute datapaths **327** (e.g., data plane development kit (DPDK) datapaths) that implement one or more LFEs. In some embodiments, SRs of logical routers are assigned to edge devices and implemented by these edge devices (the SRs are centralized, and thus not distributed in the same manner as the DRs or logical switches). The datapaths **327** of the edge devices **325** may execute in the primary operating system of a bare metal computing device and/or execute within a VM or other data compute node (that is not a logical network endpoint) operating on the edge device, in different embodiments.

(51) In some embodiments, as shown, the edge devices **325** connect the physical sites **305-315** to

each other (and to external networks **125**). In such embodiments, the host computers **320** within a physical site can send data messages directly to each other, but send data messages to host computers **320** in other physical sites via the edge devices **325**. When a source logical network endpoint (e.g., VM **330**) in the first physical site **305** sends a data message to a destination logical network endpoint (e.g., VM **335**) in the second physical site **310**, this data message is first processed by the PFE **340** executing on the same host computer **320** as the source VM **330**, then by an edge device **325** in the first physical site **305**, then an edge device **325** in the second physical site **310**, and then by the PFE **345** in the same host computer **320** as the destination logical network endpoint **335**.

(52) More specifically, when a logical network endpoint sends a data message to another logical network endpoint, the PFE executing on the host computer at which the source logical network endpoint resides performs logical network processing. In some embodiments, the source host computer PFE set (collectively referred to herein as the source PFE) performs processing for as much of the logical network as possible (referred to as first-hop logical processing). That is, the source PFE processes the data message through the logical network until either (i) the destination logical port for the data message is determined or (ii) the data message is logically forwarded to an LFE for which the source PFE cannot perform processing (e.g., an SR of a logical router).

(53) FIG. **4** conceptually illustrates an example of a network management system **400** of some embodiments for a logical network that spans physical sites **405-415**. The network management system **400** includes (i) a global manager **420** that manages the entire logical network spanning all of the physical sites **405-415** as well (ii) the local managers **425-435** for each of the sites that manage the logical network at their respective sites. Each physical site also includes central controllers, host computers, and edge devices (not shown) in addition to the local manager. In some embodiments, the global manager **420** executes on a computing device **445** at one of the sites **405** spanned by the logical network, and the local managers **425-435** also execute on computing devices **450-455** at their respective sites **410-415**.

(54) In some embodiments, the global manager **420** receives a global desired configuration for the logical network via one or more user clients **440**. Each of the local managers **425-435** also receives in some embodiments a (site-specific) desired configuration for the logical network via the user clients **440**. The desired configuration is provided to the managers **420-435** and **460** from a user client **440** in some embodiments using a representational state transfer (REST) application programming interface (API), and is represented by dashed lines in FIG. **4**. The global manager **420** also provides a site-specific portion of the global desired configuration to each of the local managers **425-435**, as represented by dotted lines in FIG. **4**.

(55) In some embodiments, as illustrated in FIG. **4**, the global manager **420** executes on the same computing device **445** at a given physical site **405** as the local manager **425** managing that site. In other embodiments, as illustrated in FIG. **5**, the global manager **420** executes on a computing device **505** at a given physical site **405** separately from any local manager **425** managing that site. In still other embodiments, as illustrated in FIG. **6**, the global manager **420** executes on a computing device **605** at a separate site **610** that is not spanned by the logical network, and therefore has no local managers.

(56) Some embodiments employ a secondary global manager **460**, in an active-standby arrangement with the (primary) global manager **420**. The primary global manager **420** is asynchronously synchronized (e.g., not real-time) with the secondary global manager **460** as a standby for failover scenarios. This asynchronous replication is represented by a dot-dash line in FIG. **4**. For resiliency, the secondary global manager **460** is located at a different physical site **410** than the site **405** where the primary global manager **420** is located, and maintains an independent database. This ensures that a failover scenario due to connectivity issues to the active global manager's physical site does not also automatically affect the secondary global manager's physical site, and exploits the multisite architecture of the federated logical network.

(57) The secondary global manager **460** executes in some embodiments on the same computing device **450** as the local manager **430** managing its site **410**, as illustrated in FIG. 4. In other embodiments, as illustrated in FIG. 5, the secondary global manager **460** executes on a different computing device **510** at its physical site **410**, separately from any local manager **430** managing that site **410**. In still other embodiments, as illustrated in FIG. 6, the secondary global manager **460** executes on a computing device **615** at a separate site **620** that is not spanned by the logical network, and therefore has no local managers. Even though FIGS. 4-6 illustrate examples where the primary global manager **420** and the secondary global manager **460** have identical hosting arrangements, in some embodiments any combination or permutation of hosting may be employed as required. As just one example, the primary global manager **420** may be co-located with a local manager (as in FIG. 4), and the secondary global manager **460** may execute at a physical site that is not spanned by the logical network (as in FIG. 6).

(58) The primary global manager **420**, the secondary global manager **460**, and the local managers **425-435** are in some embodiments separate modules of a single application, and in other embodiments are separate applications. These applications in some embodiments execute as one or more processes within machines that execute on host computers at each physical site. Some embodiments deploy one or more of the managers **420-435** and **460** as a cluster of machines at their physical site, with each machine executing on a different computing device at the same site, as described in further detail below with reference to FIG. 8.

(59) FIG. 7 conceptually illustrates different components of the network management system **400** described in FIG. 4. The desired configuration of the logical network is received by the primary global manager **420** from a user client **440** (e.g., over a wide area network **442** such as the Internet), stored in a database **710**, and replicated to the secondary global manager **460** for storage in a separate database **712**. A site-specific portion of the desired configuration is also provided to each of the local managers (via dedicated asynchronous channels) **425-435** for storage in their respective databases (not shown).

(60) The global manager **420** also includes a number of additional modules, including an API processor **706** for receiving the user client input via a REST API, a core global manager service **707** that writes data to the database **710**, a persistent work queue **711** in the database **710** to maintain causality for incoming create/update/delete (CUD) events, a log replication module **730** to replicate CUD events to the database **460** at the secondary global manager **460**, a broker/span service **735** to perform span calculations on the CUD events (and the logical network elements referenced therein), an asynchronous replication (AR) module **300** which includes dedicated persistent queues **725-727** for disseminating CUD events to different local managers at different physical sites, and site managers **745** for maintaining connection parameters used by the AR module to establish channels to the other local managers. Each of these modules and their functionality are described in further detail below.

(61) In some embodiments, the databases **710** and **712** are distributed databases (e.g., a shared log) implemented across a set of storage devices at the managers' respective physical sites. In addition, in some embodiments, the global managers **420** and **460** are implemented as a cluster of machines executing on separate computing devices at its respective physical site. FIG. 8 conceptually illustrates a manager cluster and a distributed database system at a physical site spanned by the logical network. The manager cluster illustrated in this example may function in different embodiments as a primary global manager in active mode, as a secondary global manager in standby mode, or as a local manager. The manager is an application that runs on several manager machines **805-815** executing on separate host computers **820-830** at the physical site, which communicate with each other over a physical network **835** at the physical site (e.g., a datacenter network fabric). One of the manager machines **805** is assigned as the leader for the cluster in some embodiments, which shares tasks with the other manager machines in either active/active mode (e.g., using load balancing), active/standby mode, or some combination of these modes. The cluster

of manager machines **805-815** appears as a single logical manager to other managers in the logical network.

(62) In some embodiments, the manager application that runs on the machines **805-815** has separate modules for the global manager and local manager, which can be enabled or disabled as required. Alternatively or conjunctively, in some embodiments, at least one of the machines **805-815** is a dedicated global manager machine, with a separate machine (not shown) for running an application for the local manager, executing on either the same host computers **820-830**, or different host computers (as described above with reference to FIGS. **4-6**).

(63) The manager cluster stores desired configuration data in a distributed database system that is managed by one or more instances **840-850** of a database that execute on the host computers **820-830** in some embodiments. The database executes within the local manager machine on the host in some embodiments, though they are shown as separate in the figure for clarity. The database instances **840-850** communicate with each other over the physical network **835** at the physical site (e.g., the datacenter network fabric) that is used by the manager cluster. The database instances **840-850** collectively appear to the manager cluster as the single logical database **710**. In some embodiments, the instances are shards or slices of the database. In other embodiments, each instance is a node with a full copy of the data (e.g., as illustrated in the example of FIG. **8**). The redundancy allows for durability in case one of the hosts **820-830** fails.

(64) In some embodiments, the database instances **840-850** are nodes of a distributed log that is stored on the host computers **820-830**. Entries in the distributed log provide an ordered, persisted history of updates to the state of different logical network elements and logical network policies, which the manager cluster accesses via application programming interfaces (APIs) provided by the database instances **840-850**. The distributed log and the database APIs are described in more detail by U.S. Pat. No. 10,540,119, which is incorporated herein by reference.

(65) Returning to the example of FIG. **7**, data describing the global desired configuration is received from the user client **440** and stored in the database **710** in some embodiments using a series of transactions, initiated through a series of REST API calls from the user client **440** to the primary global manager **420**. These API calls are received and processed by an API processor module **706** in some embodiments, which then provides the received data to a manager service **707** that performs the core functions of the global manager **420**. The manager service **707** stores the received data in the database **710**. In some embodiments, the data is stored in the database in tables that store configuration parameters for the logical network elements of the logical network. In some such embodiments, the data in the tables is expressed as a hierarchical tree of user intent, as described below with reference to FIG. **17**. Furthermore, in some embodiments, the manager service **707** also stores (e.g., duplicates) the incoming series of transactions in a work queue **711**, to preserve their causality and order. The work queue is a persistent queue in the database **712** in some embodiments, and for redundancy is also replicated across the database nodes.

(66) In some embodiments, the manager service **707** is the cluster of manager machines **805-815**, as described above with reference to FIG. **8**. As illustrated in FIG. **7**, in some embodiments the secondary global manager also has a corresponding API processor and manager service, which are activated in a failover scenario when the primary global manager fails (e.g., becomes unreachable) and the secondary global manager becomes the active manager.

(67) In some embodiments, the database **710** generates one or more update streams from the series of transactions. FIG. **9** conceptually illustrates generating an update stream for use by the primary global manager **420**, to replicate the desired configuration to the secondary global manager **460**. FIG. **10** illustrates a process **1000** performed in some embodiments by a database instance **840** to generate the update stream, with reference to FIG. **9**.

(68) The process **1000** begins by receiving at **1005** data describing a desired configuration of the logical network. The received data is in some embodiments one or more create, update, or delete (CUD) events received at the global manager **420** as a series of API transactions, each CUD event

affecting one or more logical network elements spanning one or more of the physical sites. For example, in FIG. 9 a CUD event **905** is received by a database instance **840** from the global manager service **707**.

(69) At **1010**, the process **1000** tags the received CUD event using metadata associated with each transaction, such as timestamp information that can be used for data ordering, and database status to prevent race conditions for access. The metadata also includes in some embodiments parameters associated with the API call, such as user ID, source address, etc. Furthermore, in some embodiments, the metadata includes span information that was specified (e.g., by an administrator) when the CUD event was defined through the user client **440**. In the example of FIG. 9, the database instance **840** appends the metadata **910** (denoted by “m”) to the data prior to storing it in the database **710**.

(70) In some embodiments, not all data stored by the global manager in the database is necessarily intended for replication. Data to be replicated in some embodiments includes policy tables, permissions, physical site information, and other data that the secondary global manager would require in order to assume active status in the event of failure of the primary global manager. Other database tables, such as those that pertain to managing the network in active mode, are not necessary for replication to the secondary global manager in standby mode. In addition, state information about the realization status of the logical network would not need to be replicated to the secondary global manager, since the realization status would be obsolete by the time a failover scenario occurred. Some embodiments distinguish data to be replicated from data that is not to be replicated, by tagging the data for replication.

(71) In some embodiments, the process **1000** also tags (at **1010**) the CUD event with an identifier that indicates that the event is to be replicated. For example, in FIG. 9 the database instance **840** appends an identifier **915** (denoted by “R”) to the event intended for replication, prior to storing it in the database **710**. The database instance **840** does not append any identifier to data that is not intended for replication in some such embodiments.

(72) At **1015**, the process **1000** stores the tagged CUD event, along with at least some of the metadata, in the database **710**. Specifically, in some embodiments, the process **1000** stores the tagged data in each of the database nodes **840-850**. In embodiments where the database **710** is a distributed shared log, the tagged data is appended to the log. In the example of FIG. 9, the tagged data is appended to position **917** of the database **710**, which is the most recent position in the database.

(73) As noted above, in some embodiments the database **710** is shared by the primary global manager **420** with a local manager **425** (e.g., on the same computing device **445**, as illustrated in the example of FIG. 4). In such embodiments, the local manager also writes data corresponding to CUD events to the database, separate from the data written by the global manager. In the example of FIG. 9, the database instance **840** also receives a CUD event **920** from a local manager service **910**, which received the event from a user client **440** via a series of API transactions. Unlike the event **905** received from the global manager service **707**, the event **920** affects one or more logical network elements that only span the physical site **405**. In this example, the database instance **840** received the event **920** from the local manager service **910** before it received the event **905** from the global manager service **707**. The database instance **840** appended metadata **925** and a different identifier **930** (denoted by “L”) to the event **920** prior to storing it in the database. Accordingly, this tagged event was written to position **932** of the database **710**, which at that time was the most recent position of the database.

(74) The process **1000** pushes at **1020** the CUD event **905** and associated metadata **910** to the log replication module **730**. In some embodiments, the process **1000** pushes the CUD event to a work queue **711** in the database (not shown in FIG. 9), from which the log replication module **730** (and, the broker service **735**, as described below) retrieves it, e.g., after a notification. The database instance **840** does not include the identifier **915** in some embodiments when pushing the CUD

event to the log replication module **730**. In some embodiments, the process **1000** also notifies the global manager that the CUD event **905** has been pushed to the queue, for dissemination to relevant local managers, as described with reference to FIGS. **12-11** below. As illustrated in FIG. **9**, data tagged with a different identifier, or not tagged with any identifier, are not pushed to the log replication module **730**. The process **1000** then ends.

(75) As noted above, the database **710** generates an update stream which pushes newly-written CUD events to a log replication module **730**, for replication to the secondary global manager **460**. Returning to FIG. **7**, the global manager **420** also includes an asynchronous replication (AR) module **740** in some embodiments, which has several persistent queues **725-727**. Each of these persistent queues **725-727** is dedicated for dissemination of desired configuration data to one of the local managers **425-435** at each physical site **405-415**. These queues are replicated across the manager cluster in some embodiments, so that in the event of failure, causality of the database events is preserved and the network management system can pick up where it left off after that failure. In some embodiments, these queues are stored in active memory (e.g., rather than on disk).

(76) When the primary global manager **420** receives the global desired configuration for the logical network, the global manager stores various portions of the global configuration in the persistent queues **725-727**, based on the relevance of the portions to the configuration of the logical network at the queue's corresponding physical site. In some embodiments, a broker service **735** of the global manager **420** identifies the relevant portions the global desired configuration for each physical site, for example based on the span of the logical network elements, as described in further detail below. Span is determined in some embodiments by a span service, which in some embodiments is part of the broker service **735** (as depicted in FIG. **7**), and in other embodiments is a standalone service (as depicted in FIG. **11** below). The desired configuration is received in some embodiments as one or more create, update, or delete (CUD) events received at the global manager **420** as a series of API transactions, with each CUD event affecting one or more logical network elements spanning one or more of the physical sites. Each CUD event is a portion in some embodiments of the global desired configuration.

(77) FIG. **11** conceptually illustrates the broker service **735** of the global manager retrieving a CUD event from the database **712** and populating the relevant persistent queues **725-727** of the AR module **740**. FIG. **12** illustrates a process **1200** performed in some embodiments by the broker service **735** to populate the persistent queues **725-727**, with reference to the example of FIG. **11**.

(78) The process **1200** begins at **1205** by receiving a notification of a new CUD event that has been stored in the database **710**. This notification is sent from the global manager **420** or a database instance **840** in some embodiments (e.g., as part of operation **1020** of process **1000**). In response to the notification, the broker service **735** retrieves the CUD event. In other embodiments, the broker service **735** periodically queries the database for new CUD events. In either case, in some embodiments, the CUD events are stored in a work queue **711** in the database **712**, to preserve their causality and order. The broker service **735** then sequentially retrieves each CUD event from this queue. In the example of FIG. **11**, the broker service **735** retrieves CUD event **1120** (labeled, "F") from the work queue **711**, after having already received CUD events **1125-1145** (labeled "A" through "F").

(79) At **1215**, the process **1200** determines whether the database **710** already stores span information **1147** for the CUD event. In some cases, the span for the CUD event is pre-defined (e.g., when specifying the CUD event through the user client **440**) and is stored in the database **712** as metadata (e.g., metadata **910**) associated with the event. If the span for the CUD event is available, at **1217** the process **1200** retrieves the span information and proceeds to operation **1230**, which is described below.

(80) If the span for the CUD event is not available, then the process **1200** determines at **1220** which physical sites are spanned by the CUD event. In some embodiments, the broker service **735** invokes a span service **736** to perform a span calculation to determine the span, taking into account

the definition of the logical network elements that are referenced by the CUD event, as well as the relationship of those logical network elements with other logical network elements in the logical network. Span calculations are described in further detail below with reference to FIG. 17. After determining the span for the CUD event, at **1225** the process **1200** also updates the database with the span information in some embodiments, by storing it as metadata associated with the event. (81) Based on the span information, at **1230** the process **1200** stores a copy of the event in each persistent queue that corresponds to one of the spanned physical sites. In the example of FIG. 11, after either retrieving or determining the span **1147** for CUD event F **1120**, the broker service **735** determines that the span is equal to site A **405** and site C **415**. The broker service **735** then copies CUD event F **1120** to the corresponding queues **725** and **727**, respectively. The process **1200** then ends.

(82) As noted above, a CUD event refers to one or more logical network elements in some embodiments. If the CUD event refers to different logical network elements with different span attributes, then in some embodiments the broker service **735** copies the relevant portions of the CUD event (corresponding to the different logical network elements) to the relevant queues instead of the entire event.

(83) As noted above, the persistent queues are first-in, first-out (FIFO) queues. In the example of FIG. 11, the broker service **735** first received CUD event A **1145** and determined the span to be sites A **405** and B **410**. Therefore, event A **1145** was copied to queues **725** and **726**. Since event A **1145** was the first event in the sequence to have been queued, it is the first event to be transmitted (as shown in FIG. 11) from the queues to the respective local managers at the respective sites. Since the primary global manager is also located at site A **405**, event A **1145** is transmitted to the local manager **425** over the physical network at the physical site (e.g., a datacenter network fabric). However, since site B **410** is a different physical site, event A **1145** must be transmitted to the local manager **430** over a wide-area network (e.g., the Internet).

(84) The subsequent CUD events B-F **1120-2940** are similarly queued in the same order as received, as illustrated in FIG. 11, with event B **1105** spanning site A **405** and C **415**, event C **1135** spanning site A **405**, event D **1130** spanning sites B **410** and C **415**, event E **1125** spanning sites A **405** and B **410**, and event F **1120** spanning sites A **405** and C **415** (as discussed above).

Accordingly, queue **1105** stores events B **1140**, C **1135**, E **1125**, and F **1120**; queue **1110** stores events D **1130** and E **1125**, and queue **1115** stores events B **1140**, D **1130**, and F **1120**. The events are transmitted from each FIFO queue in the same order as they were received.

(85) Returning to FIG. 7, in some embodiments, the AR module **740** at the global manager **420** maintains a set of asynchronous channels that connect the primary global manager **420** to the local managers **425-435**. These channels are depicted as dotted lines from the persistent queues **725-727** to corresponding AR modules **755-765** at the local managers **425-435**.

(86) In some embodiments, the AR module **740** maintains a dedicated asynchronous channel that connects the primary global manager **420** to the secondary global manager **460**. In other embodiments, replication to the secondary global manager **460** is handled by a dedicated log replication module **730**, which was described above with reference to FIG. 9. The log replication module receives the data to be replicated from the work queue **711** in the database **712** as described above, and in some embodiments replicates the data, (as depicted by a dot-dash line) directly to the database **712** of the secondary global manager **460**, bypassing the AR module.

(87) In some embodiments the secondary global manager **460** also has a receiver module **742** for establishing the connection to the log replication module **730** of the primary global manager, receiving the replicated data, and writing the data to the database **712**. If the secondary global manager **460** becomes active (e.g., due to failover), then in some embodiments the receiving module **742** assumes the role of the replication module.

(88) The AR modules **740** and **750-765** maintain the channels between the physical sites, and in some embodiments guarantee various connection parameters (e.g., the minimum bandwidth, the

maximum roundtrip time, etc.) that are required for replication of data to the secondary global manager and dissemination of data to the local managers. In embodiments where the secondary global manager and/or the local managers are implemented as a cluster of machines, the channels also identify the leader machine for each manager cluster.

(89) Some embodiments execute a set of site managers **745** at the primary global manager **420** that provide information (e.g., connectivity statistics, IP addresses, etc.) about the physical sites to the AR module **740** to use in maintaining the channels, each site manager corresponding to one of the physical sites **405-415**. The site managers **745** execute separately from the AR module **740** in some embodiments (as depicted in FIG. 7), or run as a sub-module of the AR module in other embodiments. The secondary global manager **460** also executes a set of site managers in such embodiments (not shown), though these are not active until a failover scenario.

(90) As noted above, the broker service **735** enqueues CUD events to the persistent queues **725-727** of the AR module **740**. In some embodiments, the AR module **740** polls its persistent queues **725-727** to determine if there are new CUD events. If data (e.g., corresponding to CUD events) is found in a queue, the AR module **740** retrieves the data (i.e., dequeues the CUD events) and transmits the retrieved data over the channel corresponding to that queue to the AR module of the local manager at the corresponding site. FIG. 13 conceptually illustrates a process **1300** performed in some embodiments by the sending AR module **740** at the primary global manager **420** for providing the retrieved data to one of the local managers **425-435**. The process **1300** is described below with reference to FIG. 7.

(91) The process **1300** begins at **1305** by querying a persistent queue that is dedicated to one of the local managers **425-435** at one of the physical sites **405-415**. The process **1300** queries the selected queue to see if there are any new CUD events in the queue. The queue will potentially include CUD events that were copied to the queue by the broker service **735** after performing a span calculation, as described in FIG. 12 above. In some embodiments the CUD events reference one or more logical network elements that each span one or more physical sites.

(92) If the process **1300** determines (at **1310**) that there is no new data in the selected queue, then the process **1300** returns to **1305** after a specified period of time. The period of time in some embodiments is a time-out value, which varies according to the specific connection parameters of the channel to the physical site. This partly depends in some embodiments on the location of both the global manager **420** and the specific local manager at the site associated with the persistent queue (e.g., the different scenarios for placement of the global manager **420** and the local managers, as described above in FIGS. 4-6). If the process **1300** determines (at **1310**) that there is new data in the selected queue, then the process retrieves (at **1315**) the data from the queue.

(93) At **1320**, the process **1300** then provides the retrieved data to the corresponding physical site, through a channel that connects the AR module **740** of the primary global manager **420** and the corresponding AR module of the physical site's local manager. For example, if the selected queue is queue **725**, then the process **1300** transmits the data to the AR module **755** of the local manager **425** at the same site **405**, over the network fabric of site **405** through its dedicated channel (represented by a dotted line in FIG. 7). As another example, if the selected queue is queue **726**, then the process **1300** transmits the data to the AR module **760** of the local manager **435** at physical site **415**, over wide area network **771** through its dedicated channel (represented by a dotted line in FIG. 7).

(94) In this manner the AR module **740** sends the data to the destination corresponding to the queue from which it retrieved the data, whether that destination is at the same physical site or a different physical site, through whatever intervening network infrastructure is necessary. The process **1300** then ends.

(95) FIG. 14 conceptually illustrates the architecture of a local manager receiving data through a channel from the primary global manager **420**. The local manager **430** of some embodiments includes a number of modules, including a receiving AR module **765** for receiving the data from the global manager **420**, an ingress queue **1402** of the AR module for storing received events, a

local manager service **1410** that performs the core local manager functionality, and a database **1405** (which is a distributed database in some embodiments, as described above with reference to FIG. **8**). The local manager **430** also includes in some embodiments an API processor **1415** to receive input via REST APIs from a user client **440**, and a handler **1417** to send configuration data to the control plane for realization of the desired logical network configuration. These modules and their functions are described in more detail below.

(96) In the example of FIG. **14**, the local manager **430** is at a different physical site **410** than the primary global manager's site **405**, so the data is received via a wide area network **770** (e.g., the Internet). However, the medium of communication depends in some embodiments on the location of the local manager relative to the global manager **420**. For example, if the local manager (e.g., local manager **425**) is at the same physical site as the global manager, the data is received in some embodiments through the local physical network at that site (e.g., a datacenter fabric). As another example, if the local manager is co-located on the same host machine (e.g., as illustrated in FIG. **4**), then the data is received in some embodiments through an inter-process communication method.

(97) As described above with reference to FIG. **12**, in some embodiments the broker service **735** retrieves new data from the database **710** and performs a span calculation to determine which queues **725-727** the data should be stored. The global manager's sending AR module **740** then retrieves the data from each queue, as described above with reference to FIG. **13**, and transmits the data through dedicated channels to the corresponding physical sites. Each channel connects the sending AR module **740** of the primary global manager **420** to one of the receiving AR modules **755-765** of the local managers **425-435**. FIG. **15** illustrates a process **1500** performed in some embodiments by a local manager when receiving data from the global manager, with reference to the example of FIG. **14**.

(98) The process **1500** begins at **1505** by receiving data from the primary global manager. Specifically, the data is received from the sending AR module **740** of the primary global manager, via a dedicated channel between the sending AR module and the receiving AR module of the local manager (e.g., AR module **765** of local manager **430** in FIG. **14**). In some embodiments, the receiving AR module **765** maintains a persistent ingress queue **1402** to store the received data, which in some such embodiments is stored in active memory. The received data corresponds in some embodiments to one or more CUD events that each reference one or more logical network elements that span one or more physical sites, including the physical site **410**. The order and causality of the received data is preserved by the ingress queue **1402** in some embodiments.

(99) At **1510**, the process **1500** determines if the CUD event (or events) associated with the received data is valid. The validation is based on whether there is any error or inconsistency in applying the CUD event to the configuration of the logical network at the physical site. In addition, other validations are rules in some embodiments that govern whether a logical network element can be updated to stretch its span to a new site. For example, prior to stretching an element to a new site, the security policies that affect the element must be stretched to the new site first. Otherwise, a loophole is created where the security policies applicable to the logical network element are not fully applied. In some embodiments, the validation is performed by the local manager service **1410**, which retrieves the CUD event from the ingress queue **1402** and stores it in the database **1405**.

(100) If the process **1500** determines that the CUD event is invalid, then a notification for the primary global manager of the failure to validate the CUD event is generated at **1520**. The notification in some embodiments is a notification event that is queued in an egress queue (not shown) of the AR module **765**, to be sent back to the AR module **740** at the global manager **420** via the same asynchronous channel (e.g., the dotted line in FIG. **14**). In other embodiments, the notification event is sent via an out-of-band notification channel. Notification events in the egress queue are retrieved and sent over the channel separately from process **1500**, e.g., as part of the core functionality of the local manager service **1410**, or as another module (not shown).

(101) Invalid events are not dropped in some embodiments, but are also persisted in order to

maintain causality. The invalid event will still be accepted, and an intervention (e.g., by an administrator of the network) will be required to resolve the invalidity. The notification events are described in more detail with reference to FIGS. 30-32 below. Accordingly, the process **1500** then continues to **1530**.

(102) At **1530**, the process **1500** applies the CUD event to the local desired configuration of the logical network at the physical site. The desired configuration of the logical network is expressed as a policy tree in some embodiments, which is described in further detail below with reference to FIGS. 17-20. In the example of FIG. 14, the desired configuration of the logical network, e.g. the configuration of the logical network elements whose span includes the physical site **410**, is stored in the local database **1405**.

(103) For example, if the validated CUD event is a create event, then a logical network element defined by the event is created within the desired configuration stored in the database **1405**. If the CUD event is an update event, then the desired configuration of a logical network element referenced by the event is updated within the desired configuration stored in the database **1405**. If the CUD event is a delete event, then a logical network element referenced by the event is deleted within the desired configuration stored in the database **1405**.

(104) At **1540**, the process **1500** uses the (now modified) desired configuration of the logical network to generate and provide configuration data to the control plane of the logical network (e.g., a central controller or cluster of controllers at each site). In some embodiments, these controllers identify computing devices at the site which execute physical forwarding elements, and distribute the configuration data to the identified computing devices. In some embodiments, different logical network elements span different computing devices (e.g., host computers, edge devices, etc.). Each logical network element is implemented in some embodiments by physical forwarding elements executing on the identified computing devices at the sites that are spanned by that logical network element. In other words, a logical network element is implemented by at least one physical forwarding element at each site which it spans. Some embodiments have local controllers (also referred to as chassis controllers) that execute on one or more of the computing devices alongside the physical forwarding elements, and which receive the configuration data from the controller cluster. The local controllers use the configuration data to configure the physical forwarding elements to implement the logical network elements at each computing device. The process **1500** then ends.

(105) FIG. 16 conceptually illustrates a process **1600** performed in some embodiments by a local manager at a physical site, when it receives a CUD event directly from a user client **440**, instead of from the global manager **420**. This scenario occurs for example when a local administrator of the physical site (who may or may not be the same as the administrator of the global federated logical network as a whole) modifies the logical network's desired configuration as implemented at the local site (e.g. by specifying a series of create, update, or delete events for logical network elements whose span includes the local site).

(106) The process **1600** begins at **1605** by receiving a CUD event directly from a user client **440**. For example, as illustrated in FIG. 14, data describing the CUD event is received from a user client **440** and directly stored in the database **710** in some embodiments using a series of transactions, initiated through a series of REST API calls from the user client to the primary global manager **420**. The user client **440** is not at the same physical site in some embodiments, so the CUD event is received by the local manager **430** over a wide-area network **442** (e.g., the Internet). These API calls are received and processed by an API processor module **1415** of the local manager **430** in some embodiments, which then provides the received data to the local manager service **1410** that performs the core functions of the local manager **430**.

(107) The process **1600** determines at **1610** whether the CUD event (or events) is valid. The validation is based on whether there is any error or inconsistency in applying the CUD event to the configuration of the logical network at the physical site. In some embodiments, the validation is

performed by the local manager service **1410**, either directly upon receipt from the API processor **1415**, or after retrieving the event from the database **1405**.

(108) If the process **1600** determines that the CUD event is invalid, then a notification for the failure to validate the CUD event is generated at **1617**. The notification in some embodiments is a notification event, that is provided to the user client **440** for intervention (e.g., by an administrator of the network). The process **1600** then ends.

(109) If the process determines that the CUD event is valid, then the process determines at **1615** whether the event is a local event. In other words, it determines if the CUD event only references logical network elements defined at the local site. These elements, if defined through the local manager, have no span beyond the physical site, and are not known to the global manager **420** in some embodiments. If the CUD event is a local event, then the process **1600** continues to **1635**, which is described below.

(110) If the process **1600** determines that the CUD event is not a local event, i.e. it references a logical network element that was defined at the global manager **420**, then the process determines at **1620** whether the event overrides the globally-defined desired configuration of the logical network element. This determination is made in some embodiments by applying a set of priority rules to the CUD event to determine whether the CUD event is allowed to override the globally-defined desired configuration. For example, some embodiments only allow overriding of the desired configuration by a local CUD event for networking-related configurations (e.g., message forwarding rules and policies) or configuration profiles (timers, etc. which are affected by the local site's parameters, such as latency). In such cases, the local CUD event would have priority.

(111) As another example, some embodiments prevent overrides of the desired configuration by a local CUD event for security-related configurations. In such cases, the globally-defined desired configuration would have priority. In addition, in some cases the event is an emergency-related event, which is only recognized by the local manager and therefore does override any related global configuration. If the event does not have priority to override the global configuration (e.g., according to the priority rules), then the process continues to **1617**, which was defined above.

(112) At **1625**, if the process **1600** determines that the CUD event does have priority to override the globally-defined desired configuration, then a notification for the primary global manager of the override event is generated at **1625**. The notification in some embodiments is a notification event that is queued in an egress queue (not shown) of the AR module **765** to be sent back to the AR module **740** at the global manager **420**, via the same asynchronous channel (e.g., the dotted line in FIG. **14**). In other embodiments, the notification event is sent via an out-of-band notification channel. Notification events in the egress queue are retrieved and sent over the channel separately from process **1600**, e.g., as part of the core functionality of the local manager service **1410**, or as another module (not shown).

(113) At **1630**, the process **1600** creates a local copy of the logical network element the configuration of which is to be overridden by the CUD event. The original logical network element from the global manager **420** remains as a read-only object in the local manager's database **1405**. This local copy (also referred to as a shadow object) is the target of the CUD event instead of the original.

(114) At **1635**, the process **1600** applies the CUD event to the local desired configuration of the logical network at the physical site. If the event is an override, then the process **1600** applies the CUD event to the shadow copy of the object instead of the original object received from the global manager **420**. The desired configuration of the logical network is expressed as a policy tree in some embodiments, which is described in further detail below with reference to FIGS. **17-20**. In the example of FIG. **14**, the desired configuration of the logical network, e.g. the configuration of the logical network elements the span of which includes the physical site **410**, is locally stored in the database **1405**. The CUD event is applied to the local desired configuration.

(115) For example, if the CUD event is a create event, then a logical network element defined by

the event is created within the local desired configuration stored in the database **1405**. If the validated CUD event is an update event, then the desired configuration of a logical network element referenced by the event is updated within the local desired configuration stored in the database **1405**. If the validated CUD event is a delete event, then a logical network element referenced by the event is deleted within the local desired configuration stored in the database **1405**.

(116) At **1640**, the process **1600** uses the desired configuration of the logical network to generate and provide configuration data to the control plane of the logical network (e.g., a central controller or cluster of controllers at each site). The process **1600** then ends.

(117) As noted above with reference to FIGS. **15-16**, the local manager **430** generates and provides configuration data from the desired configuration of the logical network stored in the local database **1405**. In the embodiment exemplified by FIG. **14**, the local manager service **1410** generates the configuration data from the stored desired configuration, and provides the generated data to a handler module **1417**. The handler module **1417** then distributes the configuration data to a central controller cluster **1420** of one or more controllers. The controller cluster **1420** identifies host computers **320** and edge devices **325** to which to distribute the configuration data. At each edge device **325**, the configuration data is used to configure the edge's datapath **327**. At each host computer **320**, the configuration data is received by a chassis controller **1425**, which uses it to configure at least one PFE **345** executing on the host **320**, to forward data messages to and from one or more attached VMs **1430**.

(118) The controllers **1420**, in addition to distributing configuration data to computing devices such as host computers **320** and edge devices **325**, receive physical network to logical network mapping data from the computing devices in some embodiments and share this information across datacenters. For example, these mappings are in some embodiments between logical addresses (e.g., MAC addresses of logical network endpoints i.e. VMs **1430** executing on the computing devices) and physical addresses (e.g., IP addresses of tunnel endpoints at the computing devices). The controllers **1420** retrieve and share tunnel endpoint to logical network address mapping data between the host computers **320** and edge devices **325** at the physical site **410**, and also share that data with the controllers at the other physical sites **405** and **415** so that those controllers can share this data with their own site's host computers and edge devices.

(119) In some embodiments, the computing devices at each site also execute machines alongside physical forwarding elements and local controllers. These machines include logical network endpoints, which are sources and destinations of data message traffic (e.g. computers, virtual machines, containers, etc.), and service machines, which perform services on the data traffic (e.g. firewalls, load balancers, etc.). A machine is located at a single site at a given time, but may be migrated between sites in some embodiments. These machines connect in some embodiments to the physical forwarding elements in order to exchange data messages in the network.

(120) In some embodiments, the global desired configuration of the logical network is expressed as a hierarchical tree (also referred to as a global policy tree) with nodes and connections between the nodes. The global policy tree is stored by the primary global manager **420** in its database **710**. A replicated global policy tree is also stored by the secondary global manager **460** in its database **712**. In some embodiments, the nodes represent logical network elements that span one or more sites and logical network policies that apply to those elements, and the connections represent relationships between the nodes (e.g., parent-child relationships, logical network connections, etc.). Cross-referencing between nodes is achieved by reference to a path through the tree's hierarchy (e.g., global-infra/sites/London/enforce-point/edge-cluster) which provides information about the span of each node. FIG. **17** conceptually illustrates an example of one such global policy tree **1700** of some embodiments, for a logical network that spans multiple physical sites.

(121) The logical network elements include logical forwarding elements that forward data in the logical network, e.g. logical routers, logical switches, etc. For example, in FIG. **17**, the global

policy tree root **1702** connects a single Tier-0 logical router **T0 1705**, two Tier-1 logical routers **T1A 1710** and **T1B 1715**, and two different types of network segments. These segments are an overlay network segment **1720** and a VLAN segment **1725**. The node for router **T0 1705** has a number of child nodes, including static route definitions **1730** and locale services **1735** and **1740** referencing physical sites **A 1760** and **B 1777**. In this example, the router **T0 1705** also spans site **C 1765**, but the corresponding locale services referencing site **C** are not shown in the figure for clarity. The node for router **T1A 1710** has a child node for a logical switch **1745**. The node for router **T1B 1715** has a child node for a locale service **1750** referencing physical site **A**.

(122) The locale service nodes for the **T0** router and the **T1** routers define these routers' span. For example, router **T0 1705** spans sites **A 1760**, **B 1777**, and **C 1765**, while router **T1B 1715** spans site **A 1760**. As more locale services are added to a **T0** or **T1** router, the router is stretched to the corresponding sites. Unlike router **T1B 1715**, router **T1A 1710** does not have a locale service child node, and instead has a reference (dashed line) to router **T0 1705**. Therefore, router **T1A 1710** inherits the span of router **T0 1705** (i.e., router **T1A 1710** spans sites **A 1760**, **B 1777**, and **C 1765**). Certain child nodes also inherit that span automatically in some embodiments. Accordingly, the static route definitions **1730** under the **T0** router **1705** also span sites **A 1760**, **B 1777**, and **C 1765**. The logical switch **1745** inherits the span of its parent router **T1A 1710**, which in turn derives its span from the reference to router **T0 1705**. Therefore, logical switch **1745** also spans sites **A 1760**, **B 1777**, and **C 1765**.

(123) Each node in the global policy tree **1700** has multiple attributes that define configuration parameters, some of which are defined by the user and others of which are inherited. In some embodiments, span is not the only attribute that is inherited by a child node from a parent node. For example, certain **T0** or **T1** routers that span more than one site have one of the physical sites assigned as a primary site, with the other sites being secondary sites. If such a logical router has multiple service router (SR) components, then the SR component at the primary site takes precedence for certain operations. This configuration is specified (e.g., by an administrator of the network) for the router and is not part of the configuration of the locale services under the router.

(124) The locale service nodes **1735**, **1740**, and **1750** have references (dashed lines) to edge clusters **1751** and **1752** at the respective sites **A 1760** and **B 1777**. As noted above, in this example the **T0** router **1705** also spans site **C 1765**, but the router's locale service for that site and therefore the corresponding reference to an edge cluster under the site **C** node **1765** is omitted for the sake of visual clarity. The locale service nodes are associated in some embodiments with the service routers described above with reference to FIG. 2. Edge clusters are described below with reference to site nodes. The local service nodes also have various types of child nodes in some embodiments, defining various different types of configuration information available at the respective site, including interfaces (e.g., logical ports), L2 VPNs, BGP services, and IPsec VPNs. Even though locale services are child nodes of other elements, they do not necessarily inherit the full span of those elements. A locale service node has the span of the single site in some embodiments (i.e., the site of the edge cluster node referenced by the local service node), so all child nodes only inherit the span of the single site to which the local service node refers.

(125) The logical switch **1745** is shown as a child node under router **T1A 1710**. Such logical switches, also referred to as segments, are restricted to the parent router if they are connected as child nodes (as in FIG. 17). However, in some embodiments logical switches are also directly connected to the global root **1702**. For example, overlay segment **1720** is directly connected to the global root **1702**, and has a reference (dashed line) to router **T1B 1715**. This allows the overlay segment **1720** to be moved to a different router if desired, by simply changing the reference to another logical router at the top level below global root **1702**. The overlay segment **1720** inherits the span of router **T1B 1715**, e.g. site **A**, but the overlay segment could be stretched automatically if another locale service were to be added to router **T1B 1715**.

(126) Another type of segment in some embodiments is a VLAN-backed segment. These are

defined with respect to a transport zone, which is a group of host devices at a single physical site. Therefore, the VLAN-backed segment can only span that single site where the transport zone is defined. In some embodiments, VLAN-backed segments are used as uplinks in some embodiments, to connect a logical router to an external physical router outside the logical network. In other words, the VLAN is between the Tier-0 router and the external router. Since multiple Tier-0 routers may connect to same external physical router, VLAN-backed segments are used in some embodiments to distinguish their traffic. Typically, connecting a logical Tier-0 router to physical router happens at a single physical site, since each site has its own connection to the wide-area network (e.g., the Internet) between the sites, i.e. a unique Internet Service Provider (ISP). Accordingly, VLAN-backed segments provide a way of logically isolating traffic from different Tier-0 routers to the same external router, even though the Tier-0 routers may be stretched across multiple sites and overlap in their span.

(127) In the example of FIG. 17, VLAN segment **1725** has a reference (dashed line) to an interface **1755** of the locale service **1750** under router T1B **1715**. The interface **1755** is limited to the span of the locale service **1750**, so by connecting the VLAN segment **1725** to the interface **1752**, the span of the VLAN segment is limited to only site A **1760** as required. If another locale service were to be added under router T1B **1715**, then the span of router T1B would stretch to include the new site, but the span of VLAN segment **1725** would be unchanged since its reference is to the interface **1755** of the local service **1750**.

(128) Interfaces in some embodiments are uplinks or service ports. Interfaces connect to logical switches or segments, and then logical network endpoints (such as virtual machines, data compute nodes, or other types of workloads) are attached to those logical switches and segments. These endpoints also have their own services, such as DNS, TCP, etc.

(129) In addition, the logical network elements include nodes for each physical site. For example, in FIG. 17, there are nodes for site A **1760**, site B **1777**, and site C **1765** under the global root **1702**. Each site has an enforcement point child node, under which specific resources are assigned, such as edge clusters, transport zones, etc. In the example, site A's edge cluster **1751** has incoming references from locale services **1735** attached to router T0 **1705** and from locale services **1750** attached to router T1B **1715**. The edge cluster **1752** at site B **1777** has an incoming reference from the locale services **1740** attached to router T0 **1705**. In some embodiments, edge clusters also have children corresponding to edge nodes **1753**, which actually execute the services such as firewalls, DHCP, etc.

(130) The logical network elements also include logical constructs in some embodiments, such as domains that are logical groupings of one or more sites (e.g., geographic regions), and groups of logical network endpoints that share one or more attributes (e.g., operating system, region, etc.). Domains are defined and represented as nodes in the global policy tree **1700** beneath the global root **1702**. The domains are defined in some embodiments at the global manager **420**, e.g. by an administrator of the logical network. Unlike sites, which represent a physical construct, domains are a logical construct, which serve as an envelope to group different logical entities together, e.g. for security purposes. For example, firewall policies or other policy micro-segmentation applied to the domain will automatically be applied to all groups of logical endpoints defined within the domain.

(131) In some embodiments, there are different types of domains. For example, some domains are specific to a single physical site, and are referred to as locations. This type of domain acts as the container for all site-wide and site-specific configuration and policies. In some embodiments, a location domain is automatically created for each physical site in the federated logical network, and cannot be modified by the user.

(132) Other domains are logical groups of one or more sites, and are referred to as regions. Regions are assigned to geographic regions in some embodiments. For example, in the example of FIG. 7, physical site A **405** may be in Paris, physical site B **410** in London, and physical site C **415** in New

York. These correspond to the site nodes A **1760**, B **1777**, and C **1765** in the example of FIG. 17, respectively. One region can then be defined (e.g., Europe), which includes physical sites A and B, and a different region defined (e.g., North America) which includes physical site C. This is useful for example in case there are different regulatory environments (e.g., the European Union's General Data Protection Regulation, or GDPR). Regions and locations, like all domains, are attached to global root **1702** and are not attached to other domains as child nodes. Some embodiments restrict each physical site to membership in a single location and a single region. In other words, a location may not have more than one physical site, and a physical site may not be a member of two regions. (133) In some embodiments, domains are only created as top-level nodes beneath the global root **1702**, and cannot be children of other domains or inherit span from other domains. Instead, the span of a domain is manually defined in some embodiments at the global manager (e.g., by an administrator of the logical network) as the sites that are members of the domain. The span is represented in some embodiments by a domain enforcement point, which is configured to reference the site enforcement point for whichever sites the domain is intended to span. These domain enforcement points are only used by the global manager in some embodiments, and are not used by the local managers. For example, in FIG. 17, the domain enforcement point for domain A **1770** references the site enforcement point of site A **1760** (e.g., Paris) and the site enforcement point of site B **1777** (e.g., London). Therefore, the domain A **1770** is a region (e.g., Europe) spanning sites A and B, as well as other sites (e.g., Berlin) that are not shown in FIG. 17. In addition, the domain enforcement point for domain B **1775** references the site enforcement point of site C **1765** (e.g., New York). Therefore, the domain B **1775** spans site C **1765**. In this example, domain B **1775** is a region (e.g., North America) that may also span other physical sites (e.g., Chicago, Los Angeles, etc.) that are not shown in FIG. 17. Alternatively, domain B **1775** is a location that is specific to site C **1765** alone. For a given domain, the group of (one or more) site enforcement points that are referenced by the domain's enforcement point is also referred to as a domain deployment map in some embodiments.

(134) In some embodiments, logical network endpoints at each site are logically organized into security groups which can span multiple sites. Service machines as well as managed forwarding elements executing on host computer apply logical network policies (such as network policy **1773**) to the data messages exchanged between security groups of endpoints in some embodiments, based on policy rules that are defined in terms of these groups. Such security groups and network policies are defined at the global manager **420** through the user client **440** (e.g., by an administrator of the logical network). In some embodiments, security groups and network policies are represented in the global policy tree **1700** as child nodes of domains, and accordingly inherit their parent domain's span. In some embodiments, the span of a network policy is defined not only by its parent domain, but also by sites and/or domains which are referenced by the policy.

(135) For example, in FIG. 17, domain A **1770** has a child node corresponding to security group A **1771**, which accordingly inherits a span of sites A **1760** and B **1777** (i.e., the span defined by the domain deployment map of domain A). In addition, domain B **1775** has a child node corresponding to security group B **1772**, which accordingly inherits a span of site C **1765** (i.e., the span defined by the domain deployment map of domain B **1775**).

(136) Domain A **1770** also has a child node corresponding to a network policy **1773**. The network policy is applicable to any groups defined under the same domain (e.g., group A **1771**). In order to apply a policy to a security group, the span of the security group in some embodiments must include the span of the policy.

(137) In some embodiments, network policies may also refer to security groups that are not in the same domain. For example, the network policy **1773** also references security group B **1772**, which is in domain B **1775**, even though the domain deployment map for the parent domain A **1770** does not include domain B **1775**. Such cross-domain policies are discussed in further detail below with reference to FIG. 25.

(138) In some embodiments, some nodes derive their span from span leaders. Certain types of nodes are span leaders, for example a T0 router, a T1 router that is below the global root, a domain, or locale services. A T1 router that is defined beneath a T0 router is not a span leader, in some embodiments. A domain is the span leader for all security policies defined under the domain. Groups are also span leaders in some embodiments, and do not strictly inherit span from their parent nodes (e.g., domains) since the span of a group may be extended beyond its parent domain through reference groups, which are described in more detail below with reference to FIG. 25.

(139) FIG. 18 conceptually illustrates five examples of security groups that span one or more physical sites. The figure illustrates a group of host computers **1805** located at physical site A (e.g., Paris), a group of host computers **1810** located at physical site B (e.g., London), and a group of host computers **1815** located at physical site C (e.g., New York). The host computers **1805** at site A execute a group of virtual machines (VMA-VIVID) each of which either runs operating system OS1 (e.g., Windows Server by Microsoft Corporation) or operating system OS2 (e.g., Linux). Likewise, the host computers **1810** at site B execute VMs VME-VMH, and the host computers **1815** at site C execute VMs VMI-VML.

(140) As noted in the discussion above with reference to FIG. 17, groups are also represented in some embodiments as child nodes of a domain. In such cases, the group inherits the span of its parent domain, i.e. the group is defined at the sites which are members of its parent domain. Logical network endpoints located at a site may join a group if the group is defined at that site (i.e., if the group spans that site).

(141) As described above with reference to FIG. 17, sites A and B (e.g., Paris and London) are members of domain A **1770**, which is a region (e.g., Europe). This logical grouping of the two physical sites is represented by a dashed line enclosing the host computers **1805** and the host computers **1810**. In addition, site C (e.g., New York) is a member of domain B **1775**, which is also a region (e.g., North America). This logical grouping is represented by a dashed line enclosing the host computers **1815**.

(142) The first security group is group A **1817**, which is defined as all VMs which run operating system OS1 in Paris. This security group is defined logically at the global manager **420** (e.g., by an administrator of the logical network) to have a span of domain A **1770** (i.e., sites A and B). As a result, VMs A and D are automatically added to security group A **1817**. Security policies which span domain A **1770** may be applied by an administrator of the network to security group A **1817**, and would automatically therefore be applied to VMs A and D.

(143) The second security group is group B **1819**, which is defined as all VMs which run operating system OS2 in New York. This security group is defined logically at the global manager **420** (e.g., by an administrator of the logical network) to have a span of domain B **1775** (i.e., site C). As a result, VMs J and K are automatically added to security group B **1819**. Security policies which span site C (i.e., domain B **1775**) may be applied by an administrator of the network to security group B **1819**, and would automatically therefore be applied to VMs J and K.

(144) The third security group is group C **1820**, which is defined as all VMs which run operating system OS1 in Europe. This security group is defined logically at the global manager **420** (e.g., by an administrator of the logical network) to have a span of domain A **1770** (i.e., sites A and B). As a result, VMs A, D, E, and G are automatically added to security group C **1820**. Security policies which span site A, site B, or sites A and B may be applied by an administrator of the network to security group C **1820**, and would automatically therefore be applied to VMs A, D, E, and G.

(145) The fourth security group is group D **1825**, which is defined as all VMs which run operating system OS2 in Europe. This security group is defined logically at the global manager **420** (e.g., by an administrator of the logical network) to have a span of sites A and B. As a result, VMs B, C, F, and H are automatically added to security group D **1825**. Security policies which either span site A, site B, or sites A and B may be applied by an administrator of the network to security group D **1825**, and would automatically therefore be applied to VMs B, C, F, and H.

(146) The fifth security group is group E **1830**, which is defined as all VMs which run operating system OS1 in all regions globally. This security group is defined logically at the global manager **420** (e.g., by an administrator of the logical network) to have a span of sites A, B, and C. As a result, VMs A, D, E, G, I, and L are automatically added to security group E **1830**. Security policies which either span site A, site B, site C, or any combination of sites A-C, may be applied by an administrator of the network to security group E **1830**, and would automatically therefore be applied to VMs A, D, E, G, I, and L.

(147) As noted above, in some embodiments for a network policy to be applicable to a security group, the span of the security group must include the span of the network policy. For example, network policy **1773** in the example of FIG. **17** would be applicable to security group E **1830**, since the span of security group E **1830** includes the span of network policy **1773**. The applicability of network policies to security groups is discussed in further detail with reference to FIG. **25** below. In addition, in some embodiments, VMs may migrate between physical sites, e.g. VM C migrating (dashed arrow) from host machines **1805** to host machines **1810**.

(148) As discussed above, the global manager **420** performs a span calculation in some embodiments for each logical network element, to determine which logical network elements are implemented at each physical site. This is performed to provide the relevant desired global configurations to each physical site (e.g., through the asynchronous replicator channel, as described above in some embodiments with reference to FIG. **13**).

(149) FIG. **19** conceptually illustrates a process **1900** performed in some embodiments by the global manager **420** to parse the global desired configuration in order to populate the global policy tree **1700** and identify, for each physical site, a relevant portion of the global desired configuration. The process begins at **1905** by receiving the global desired configuration of the logical network from a user client **440** (e.g., from an administrator of the logical network). In some embodiments, the received configuration is a complete configuration, e.g. an initialization of the logical network. In other embodiments, the received configuration is a modification to a previously received global configuration, such as a create, update, or delete event to one or more logical network elements.

(150) At **1910**, the process **1900** populates a global policy tree (e.g., the global policy tree **1700** in FIG. **17**). based on the received global desired configuration. For example, the process **1900** identifies each logical network element referenced in the desired configuration, and creates a node in the global policy tree corresponding to the logical network element. In addition, in some embodiments each logical network element is associated with configuration parameters in the received desired configuration. These associated parameters are used to define the attributes of the created node. In some embodiments, the associated parameters in the desired configuration include references to other logical network elements. These references are used to define references between the created node and other nodes. In some embodiments, the nodes are defined in hierarchical fashion, starting from the global root **1702**, and iterating over each branch of the policy tree from the top-level nodes (e.g., logical routers, logical switches and segments, domains, sites, etc.). Some embodiments create nodes for sites first, based on auto-discovery of the physical sites, either prior to or after receiving the global desired configuration.

(151) At **1915**, the process **1900** selects a physical site that implements the logical network. In some embodiments, the physical sites are identified through auto-discovery of the physical sites as discussed above. For the selected physical site, the process **1900** parses at **1920** the global policy tree, to identify policy nodes whose span include the selected physical site. The process **1900** uses the identified nodes to generate (at **1925**) relevant desired configuration for the selected physical site. In some embodiments, the process generates the relevant configuration by identifying portions of the global desired configuration. The process **1900** provides the generated relevant configuration (e.g., the portions of the global desired configuration) to the local manager at the selected physical site, through the asynchronous replicator channel as described above in some embodiments with reference to FIG. **13**.

(152) The process **1900** determines, at **1935**, whether there are additional physical sites. If there are additional sites, the process **1900** returns to **1915**, which was described above. If there are no additional sites, the process **1900** ends.

(153) The relevant portion of the global desired configuration that is provided to each local manager in some embodiments is a site-specific policy subtree of the global policy tree, corresponding to the logical network elements implemented at that site. The subtree for each site only includes nodes for logical network elements that span the site. The global policy tree **1700** is stored by the global manager **420** in its database **710**.

(154) In some embodiments, the global policy tree is parsed to identify the span of each node in order to generate a policy subtree for each physical site. In some embodiments, the subtree is generated by first identifying the span of each node in the global policy tree. The global policy tree is then parsed, using the identified span for each node, to generate the policy subtree for each site. The local manager at each site (or a management plane application, which may be separate from the local manager) uses the relevant portion of the global desired configuration, received from the global manager, along with any desired configuration received directly by the local manager itself, to manage the logical network at the site. FIG. **20** conceptually illustrates a process **2000** performed in some embodiments by the global manager **420** to generate policy subtrees. FIG. **21** conceptually illustrates an example of a policy subtree for a single site generated by process **2000**.

(155) The process **2000** begins at **2005** by receiving a global policy tree (e.g., the global policy tree **1700** in FIG. **17**) that represents the desired configuration of the multisite logical network. The process **2000** selects at **2010** one of the nodes of the global tree, corresponding in some embodiments to one of the logical network elements. In some embodiments, the nodes are selected in an iterative fashion over the hierarchy of the global tree, starting with the top-level nodes under the global root **1702** (e.g., logical routers, logical switches and segments, domains, etc.).

(156) The process **2000** determines (at **2015**) the span attribute for the selected node. In some embodiments, the span attribute is determined by a span calculation performed by the broker service **735** of the global manager **420**. In some embodiments, the span calculation is based on the relationships between the nodes in the global policy tree. For example, when the relationship between two nodes is a parent-child relationship, the child node may inherit its span from the parent node. In other cases, however, a child node does not inherit the entire span of its parent node. As another example, when the relationship between two nodes is a dependence of one node on another node, expressed as a reference from one node to the other, the span of the dependent node will depend on the referenced node. Some logical network elements also have a pre-defined span in some embodiments, e.g. defined by an administrator of the network, which is stored in the database **710**. In such embodiments, the process **2000** retrieves the span attributes from the database.

(157) At **2020**, the process **2000** uses the span attribute to assign the selected node to a policy subtree for each physical site that is spanned by the node. These policy subtrees are stored in some embodiments in the database **710**. In some embodiments, the policy subtrees are associated with the dedicated persistent queues **725-727** that each correspond to the physical sites.

(158) The process **2000** determines (at **2025**) if there are additional nodes in the global policy tree. If there are additional nodes, then the process returns to **2010**, which was described above. If there are no additional nodes, then the process continues to **2030**. At **2030**, the process uses the generated policy subtrees to generate relevant desired configuration for each physical site. In some embodiments, the process generates the relevant configuration from the policy subtrees by identifying portions of the global desired configuration. The process **2000** then ends.

(159) FIG. **21** illustrates a global policy subtree **2100** for the physical site A **405**, stored at the database for local manager for site A (i.e., local manager **425**). Since the subtree **2100** is specific to site A, all top-level nodes below global root **2102** with span that includes site A are preserved, while top-level nodes that are relevant to sites B and C are omitted. For example, the top-level

nodes for sites B **1777** and C **1765** are omitted, as are all their respective child nodes. However, the top-level node for domain B **1775** remains, since in some embodiments all local managers are informed of all available domains. In addition, for router T0 **1705**, the locale services node for site B **1740** is also omitted. Router T0 **1705** still spans sites A and B in some embodiments, since its span attribute is associated with its definition at the global manager **420**, not the local manager **425**. (160) Network policy **1773** is also preserved in the global policy subtree **2100**. This policy **1773** is defined under domain A **1770**, so in some embodiments it has a span of site A and site B, even though this subtree is specific to site A. In addition, as noted above with reference to FIG. **21**, network policy **1773** also references group B **1772**. As a result, the span of the policy **1773** also includes site C, even though that site is in a different domain. This reference to group B **1772** is also preserved in the policy subtree **2100**, and is discussed in more detail below with reference to FIG. **25**.

(161) In some embodiments, a local manager also stores a separate policy tree, that is generated based on desired configuration received directly at the local manager instead of from the global manager **420**. FIG. **22** conceptually illustrates a local policy tree **2200** for site A that is generated in some embodiments by one or more operations of process **1900**. This local policy tree **2200** is distinct from the global policy subtree **2100** received from the global manager **420**. In this example, the local root **2202** of the local policy tree **2200** connects a site node for site A **0560**, two Tier-1 logical routers T1C **2205** and T1D **2210**, as well as an overlay segment **2215**. The overlay segment **2215** includes a reference to the router T1D **2210**, which has a locale service corresponding to site A.

(162) In some embodiments, logical network elements defined in the local policy tree **2200** may reference logical network elements defined in the global policy tree **1700**. For example, in FIG. **22**, the router T1C **2205**, references the router T0 **1705** that was defined from the global manager **420**. As a result, the router T1C **2205** is able to utilize the static routes **1730** associated with the T0 router **1705**, even though these routes are not defined locally at site A.

(163) The global policy tree **1700** is stored by the primary global manager **420** in its database **710**. A replica of the global policy tree **1700** is also stored by the secondary global manager **460** in its database **712**. As noted above, in some embodiments the nodes also represent logical network policies that apply to the logical network elements. The logical network policies include forwarding policies, service policies, and security policies, and are applied in some embodiments to govern the behavior of the logical forwarding elements (e.g., by governing the behavior of the physical forwarding elements that implement the logical forwarding elements).

(164) Policies are defined in some embodiments at the global manager **420** through a user client **440**, e.g. by an administrator of the logical network. In some embodiments, policies are one or more service rules which are enforced at the sites on data message flows based on a set of flow attributes. The global manager **420** in some embodiments distributes the service rules to local managers **425-435** at the sites **405-415** at which the policy is to be enforced (i.e., the policy's span). Each local manager uses the received service rules to generate configuration data for distribution by controllers **1420-1425**, to configure the data plane (i.e., the physical forwarding elements **345** and edge devices **325**) to enforce the received service rules on data message flows between logical network endpoints, such VMs and other data compute nodes.

(165) The policies are defined in some embodiments by reference to groups of logical network endpoints that span one or more sites, e.g. the security groups **1817-1830** described above with reference to FIG. **18**, which are defined in some embodiments at the global manager **420** (e.g. by an administrator of the logical network). The service rules refer to these groups in some embodiments by using a group identifier that is assigned at the global manager **420** when the groups are defined. The definitions of these groups are distributed to the sites **405-415** spanned by the policies, so that the controllers **1420-1425** at these sites are able to identify hosts for member VMs for each group and enforce the policy at the identified hosts.

(166) As discussed above, the service rules for a given policy are enforced on data messages that are associated with a group of logical network endpoints (e.g., VMs) spanning one or more sites. Service rules have multiple fields in some embodiments, including source fields, destination fields, and action fields. Some embodiments refer to these groups in these fields by using group identifiers that are assigned to the groups (e.g., by the global manager when the groups are defined). For example, when the group identifier for the group is referenced in a source field of the service rule, the service rule is enforced on data messages that originate from machines in that group. When the group identifier for the group is referenced in a destination field of the service rule, the service rule is enforced on data messages that are directed to machines in that group. In addition to group identifiers, the service rule is also enforced in some embodiments on data message flows with other attributes matching those specified in the rule. For example, the rule may specify a flow 5-tuple (source IP address and port number, destination IP address and port number, and transmission protocol). Alternatively or conjunctively, the rule may specify other attributes of the flow or a referenced group identifier.

(167) In some embodiments, the controllers **1420-1425** provide the service rules to identified hosts executing PFEs, and the PFEs store the service rules as match-action tables. FIG. **23** conceptually illustrates a match-action table **2300** of some embodiments. In this example, fields for source group identifier, destination group identifier, and action are shown. Other attribute fields, used in some embodiments for matching to data message flows as discussed above, are not shown for the sake of clarity.

(168) The match-action table **2300** has multiple flow entries **2305-2315** each specifying different service rules. The number of entries shown is simplified for illustrative purposes only and in practice there are hundreds, even thousands of such entries. These entries refer to the security group definitions introduced earlier with reference to FIG. **18**. For example, the first entry **2305** indicates that matching data messages originating from group C (OS1 machines in Europe) and destined for group D (OS2 machines in Europe) should be dropped. The rationale for such a rule in some embodiments may be differing security criteria between the two different operating systems.

(169) The second entry **2310** indicates that matching data messages originating from group A (OS1 machines in Paris) and destined for group E (OS1 machines globally) should be allowed. This rule indicates that OS1 machines at one site in a domain are allowed to send data messages to machines at other sites in the domain.

(170) The third entry **2315** indicates that matching data messages originating from group B (OS1 machines in New York) and destined for group A (OS1 machines in Paris) should be dropped. This rule indicates that OS1 machines in one region are not allowed to send data messages to OS2 machines in another region. In this example, the rule specified by this entry **2315** is one possible rule for network policy **1773**, which as illustrated in FIG. **17** is defined under domain A **1770**, alongside group A **1771**, but has a reference to group B **1772** defined under domain B **1775**. In this example, the policy tree reference points from the policy **1773** to group B **1772**, even though group B is the source group and not the destination group for the policy. This is because the policy tree reference is related to the policy's span, which is independent of whether the groups are source or destinations of data message traffic within that span.

(171) The PFEs identify matching flows by performing match classification operations on each data message, to determine if the flow attributes and/or group identifier match those specified in the service rule. When a data message matches the service rule, it is provided to a service machine in some embodiments, which performs the action that is specified by the rule on the matching data message. These actions include dropping the data message, allowing the data message, or performing other service operations on the data message. For example, a service operation may modify the header of the data message, to redirect the message to a different destination. These service operations include load-balancing, firewall, Dynamic Host Configuration Protocol (DHCP), Network Address Translation (NAT), and other services.

(172) In some embodiments, policies are defined as nodes in the policy tree, and defined as child nodes of domain nodes. In such cases, the policy (and therefore, any service rule of the policy) inherits the span of its parent domain, i.e. the policy is enforced at the sites which are members of its parent domain. Furthermore, groups are also represented in some embodiments as child nodes of a domain, and the group inherits the span of its parent domain. For a policy to be applicable to members of a group, the span of the group must in some embodiments include the span of the policy. This is so the controllers are able to recognize the identifiers referenced by the service rules, and identify the members of those groups to enforce the service rules.

(173) In some embodiments, a policy's service rules distributed to a first set of sites (e.g., a first domain) for applying to a first group of machines defined at those sites make reference to a second group of machines defined at a second, different set of sites (e.g., a second domain). In other words, the span of the referenced second group does not include the span of the policy. Since such service rules reference a group of machines that is undefined for one or more of the sites in the first set of sites (e.g., the first domain), not all the machines in the first group are able to enforce the policy. Therefore, some embodiments identify which sites in the first set of sites (e.g., the first domain) do not have the definition of the referenced group, and distribute the definition of the referenced group to those identified sites.

(174) FIG. 24 conceptually illustrates a process 2400 performed in some embodiments by the global manager 420 to extend the span of a group for applying a cross-domain policy. FIG. 25 conceptually illustrates an example of extending the group span due to a cross-domain policy. The policy in this example is the same policy 1773 described above in FIG. 17, and defined with respect to the security groups described above in FIG. 18.

(175) FIG. 25A conceptually illustrates simplified versions 2500 of the global tree 1700, as defined at the global manager 420, the local manager 425 at site A 1760 (e.g., Paris), and the local manager 435 at site C 1765 (New York), prior to defining any network policies. The Paris site is a member of the Europe region (domain A 1770) and the New York site is a member of the North America region (domain B 1775). As described above in FIG. 18, the Europe domain also spans site B (London).

(176) Two groups are defined beneath these regions, group A 1817 (OS1 machines in Paris) under domain A 1770 (Europe), and group B 1819 (OS2 machines in New York) under domain B 1775 (North America). Both domains A 1770 and B 1775 and groups A 1817 and B 1819 are defined at the global manager 420. Group A 1817 is defined at the local manager 425 in Paris, and group B 1819 is defined at the local manager 435 in New York. However, both domains 1770-1775 are defined at both sites, since domains are in some embodiments defined at all sites. The corresponding site nodes under each domain are omitted from the trees in FIG. 25 for clarity.

(177) The process 2400 begins at 2405 by receiving a network policy defined in a domain (e.g., which spans a first set of sites) for enforcement on members of a security group in that domain. The network policy is received at the global manager 420 through a user client 440 in some embodiments (e.g., from an administrator of the network). For example, in FIG. 25B, the security policy R 1773 is now defined at the global manager 420 as a policy beneath domain A 1770 (Europe), for enforcement on members of group A 1817. The policy 1773 is also distributed to the local manager 425 in Paris. Though not shown, the policy 1773 is also distributed to the other local managers at the other sites in domain A 1770 (Europe) (e.g. London). In the exemplified embodiment, the policy 1773 includes service rule 2315, which is distributed to the local managers in the domain.

(178) At 2410, the process 2400 identifies a reference in the network policy to another security group in another domain. For example, in FIG. 25B, the policy R 1773 defines a service rule with a source group of group B 1819 and a destination group of group A 1817. Since group B 1819 is not in the same domain as the policy 1773, the reference to group B 1819 is indicated by a dashed line in the policy tree at the global manager 420, from policy R 1773 to group B 1819. The action

specified by policy R **1773** is a drop action. In other words, the policy R **1773** specifies that incoming data message traffic is not allowed from machines in New York running OS2 to machines in Paris running OS1. Such traffic is to be dropped in Paris.

(179) The process **2400** determines at **2415** whether the span of the other group includes the span of the policy. If the group span includes the policy span, then the process **2400** ends. If the group span does not include the policy span, as in the example of FIG. **25B**, then the process continues to **2420**, which is described below. For example, in FIG. **25B**, the span of the policy is the first domain A **1770**, e.g. Paris and London. The span of the referenced group B **1819** is the domain B **1775**, e.g. New York. Since the policy references a group that is not in its parent domain, the span of the referenced group does not include the span of the policy.

(180) The process **2400** extends at **2420** the span of the referenced group include the span of the policy. In effect, this distributes the definition of the group in the other domain to the sites in the original domain where the policy is defined. This is done in some embodiments by performing a new span calculation for the policy, based on separate span calculations for the groups referenced by the policy as well as the policy's parent domain. The span of the referenced group is then extended by sending an update event through a dedicated AR channel from the global manager **420** to the local manager **425**, as described above with reference to FIG. **13** above.

(181) As noted above, in some embodiments span is also determined according to span leaders for some nodes of the global policy tree **1700**. Security groups are one type of node in some embodiments whose span leader is not their containing domain. For example, some groups serve as span leaders for other groups. In the case of cross-domain policies, the span of the referenced group is extended by specifying the span leader of the referenced group to be the first group in the same domain as the policy. In this manner, the span for the referenced group is broader than just its parent domain.

(182) Upon receiving the update event, the local manager **425** in some embodiments creates a reference group underneath the global policy tree in Paris. For example, in FIG. **25B**, a reference group **2550** is defined at the local manager **425** in Paris by the global manager **425**. In some embodiments, this reference group **2550** is not a full group definition like group A **1817**, but instead is a symbolic path reference to group B **1819**. As a result of this reference, for the local manager **425**, group B **1819** now effectively spans New York, London, and Paris (i.e., spans domains A and B). The process **2400** then ends.

(183) Some embodiments extend the group definition using process **2400** for all local managers at all sites (e.g., London) in domain A (Europe). For example, if the group definition changes, i.e. a VM migrates from one site in the domain to another site in the domain, then the destination site will also need the reference group in its local manager's global tree.

(184) In some embodiments, the policy tree at the local manager **435** in New York is not updated or modified in any way as a result of extending the span of the group **1819** at the global manager **420** and adding the reference group **2550** to the local manager **425**. The administrator of the New York site is not aware in some embodiments that the group **1819** at the site is being referenced by a policy **1773** at a different domain **1770** and could even decide to delete the group. However, due to the reference group **2550** created at the referring domain **1770** the global manager **420** is aware of the reference. The global manager **420** then prevents the deletion of the group **1819** and in some embodiments displays an alarm to the administrators of both the referring site and the referred site, so that the conflict can be resolved.

(185) After the local managers in the original domain receive the reference group, they generate configuration data and distribute to their respective site controllers, to configure the physical forwarding elements and service machines at the first set of sites to enforce the policy.

(186) In some embodiments, the configuration of the logical network can be described in terms of several different types of state. The desired state represents the user intent, e.g. the state of the network that the user (e.g., an administrator) desires to implement, through the global manager **420**.

The realized state is the actual implementation of the desired state, which includes configuration of the physical forwarding elements in the data plane. The realized state only pertains to logical network elements created by the user, and includes configuration details that the network management and control system automatically determines on behalf of the user.

(187) The third type of state is the discovered state, which represents configuration details of logical network elements that are not specified as part of the user intent. For example, one category of discovered state in some embodiments is site-specific information, such as the available physical resources at a physical site. In some embodiments, the desired state is dependent on the discovered state, or is informed by the discovered state. As the desired state is realized, some parts of the discovered state are also modified in some embodiments, and these modifications are then discovered again, informing the next iteration of desired state and realized state. The discovered state automatically discovered is in some embodiments by the global manager **420** based on information from the physical sites, communicated up from those sites' control planes to their local managers, and from the local managers to the global manager.

(188) For example, the structure of the global policy tree **1700** beneath each site node is automatically discovered in some embodiments, from the controller cluster **1420** at the physical site, up to the local manager **430**, and then to the global manager **420**. Once the global manager **420** is configured with the access credentials of the local manager **430**, it communicates with the local manager (e.g., through a REST API in some embodiments) to automatically populate the enforcement point node and identify which transport zones, edge clusters, and edge nodes are configured and available. In some embodiments, these were previously configured by the physical site's local manager. Corresponding child nodes are accordingly created for each.

(189) In some embodiments, the creation of logical ports and segment ports is not allowed from the global manager **420**. Such ports are site-specific in such embodiments, with unique identifiers that are assigned by the management/control plane at their home physical site, and the span of these ports cannot be stretched. Only one logical network endpoint can be attached to a port and a logical network endpoint can only be part of one site at a given point in time. When the span of a switch/segment gets stretched to multiple sites, active logical ports (e.g., with workloads attached) are automatically discovered as child nodes of the logical switch or segment in that tree (e.g., as interface nodes).

(190) The desired configuration of the logical network represents the intentions of the user (e.g., the administrator). The user specifies their intent by specifying the desired configuration, which is why the desired configuration is also referred to as user intent. The global manager **420** is an intent-based policy manager that receives user intent (internally represented as the global policy tree **1700**) and communicates that intent to the local managers at each site. The local managers then interpret the received user intent to generate configuration data, and provide the configuration data to the network managers and controllers as described above to implement the desired configuration. In some embodiments, this is referred to as realization of the user intent (e.g., in the control plane and the data plane). The actual implementation details (e.g., configuration parameters of the forwarding elements, IP addresses, internal UUID identifiers for ports, etc.) are not known or dictated by the user or by the global manager when specifying the intent, but are necessary details for realization of that intent. These details include data from the management plane and the control plane, e.g. site managers and controllers.

(191) In a best-case idealized scenario, the realized state of the logical network would be identical at every physical site. However, since each physical site is independently configurable (e.g., via the local managers **425-435**), and since certain configuration details are site-specific (e.g., unique identifiers for system-created objects in the management and control planes, such as logical port identifiers), in practice the realized state will be different for every physical site, even for the same logical network element. Even for a single site, the amount of realization data is vastly larger than the amount of data required to specify user intent, and this is then multiplied by the number of sites

in the federated logical network. Accordingly, in some embodiments the global manager does not automatically maintain information on realization status (also referred to as the “realized state”) for any given logical network element. Instead, the realized state may be queried from the global manager on an on-demand basis, for troubleshooting or other investigatory purposes.

(192) The realized state of a logical network element indicates in some embodiments whether the desired configuration (also referred to as the desired state) of the logical network element has been realized at each of the sites that the element spans. Some embodiments describe the realized state partly in terms of a deviation from the desired state. In some embodiments, the realized state for a logical network element (e.g., a logical router, logical switch, etc.) describes how the logical network element is implemented at each of the physical sites that it spans. For example, the realized state includes data from PFEs executing on computing devices at each of the spanned sites to implement the logical network element.

(193) For each logical network element, the realized state received in some embodiments from each local manager is a subset of the total realized state of that logical network element, just as the desired state for the logical network element that is provided to the local manager is a subset of the global desired state of that logical network element defined at the global manager.

(194) In some embodiments, the realized states received from each local manager **425-435** are status messages that indicate whether the logical network element was successfully realized at the corresponding physical site. For example, if the logical network element has not been successfully realized at all physical sites, the status message may simply be a flag that indicates error or a pending realized state.

(195) FIG. **26** conceptually illustrates a process **2600** performed in some embodiments by the global manager to retrieve the realized state of a logical network element. FIG. **27** conceptually illustrates a scatter-and-gather operation of the global manager to retrieve, on demand, the realized state of a logical network element that spans multiple physical sites. FIGS. **28-29** conceptually illustrate examples of a user interface of some embodiments for presenting the realized state to an administrator of the logical network. In some embodiments, the interfaces described in FIGS. **28-29** are a single interface presented on a user client **440**, which can select between modes of operation for accessing the global manager **420** or any of the local managers **425-435**.

(196) The process **2600** begins at **2605** by receiving an inquiry for the realized state of a logical network element (e.g., a logical router, a logical switch, services such as BGP or VPN, security policies, etc.). In some embodiments, the global manager receives the inquiry from a user client **440** (e.g., over a wide area network **442** such as the Internet), and is initiated in some embodiments by an administrator of the logical network. The inquiry is received and processed in some embodiments as one or more REST API calls from the user client **440** to an API processor module **706**.

(197) In some embodiments, the inquiry includes a reference to the logical network element, e.g. a path identifier that uniquely identifies the element. For example, an inquiry for the realization status of the logical switch **1745** includes a path identifier such as `/global-infra/router_T1A/` which describes where the node corresponding to the logical switch is positioned in the global policy tree **1700**.

(198) In some embodiments, the inquiry is for the realized state of the logical network element at a single physical site, instead of at all sites spanned by the element. In such embodiments, the inquiry includes an identifier for the specified site. The process **2600** determines at **2607** whether the inquiry includes a site identifier. If the process **2600** determines that the inquiry includes a site identifier, then the process provides the site identifier to a realization status processor **2705** and continues to **2615**, which is described below.

(199) If the process **2600** determines at **2607** that the inquiry does not include a site identifier, the process **2600** identifies at **2610** the physical sites spanned by the logical network element. Identifying the spanned sites is achieved in some embodiments by performing a span calculation as

described above with reference to FIG. 12. Specifically, in some embodiments, the API processor **706** provides parameters associated with the inquiry to the span service **736** to perform the span calculation. For example, in some embodiments the parameters include the path identifier for the logical network element. In some embodiments, the span is determined from the definition of the logical network element (e.g., the span was specified by an administrator of the logical network). After determining the span, the span service **736** provides site identifiers corresponding to the span of the logical network element to the realization status processor **2705**.

(200) The process **2600** sends at **2615** a request for the realized state of the logical network element to each local manager **425-435** at each identified site. In some embodiments, the request is sent from the realization status processor **2705** that receives one or more site identifiers from the broker service **735** or the API processor **706**. Some embodiments send the requests in parallel (e.g., “scatter” the requests) to the local managers **425-435**.

(201) In some embodiments, the requests are made through an API of the local managers **425-435**. In such embodiments, the realization status processor **2705** employs multiple realization API clients **2710** to perform the API calls, with one client assigned to each local manager. In some embodiments, the requests are transmitted through the same AR channels **725-727** as are used to provide the desired configuration data to each local manager.

(202) The process **2600** receives at **2615** realized state from each of the local managers in response to the requests (e.g., “gathers” the realized state). After receiving the realized state from the local managers at each of the sites spanned by each of the logical network elements, the process **2600** combines at **2620** the received realized state from each site. In some embodiments, the received realized states are received by the realization API clients **2710** and combined by the realization status processor **2705** in a single realization state for the logical network element.

(203) At **2625**, the process **2600** processes the single combined realization state and presents the realized state in response to the initial request. For example, the realized state may be presented in a report on a user client **440** for an administrator of the logical network in some embodiments. The process **2600** then ends.

(204) FIGS. **28A-28C** conceptually illustrate an example of a user interface **2800** for querying the realization status of logical network routers from the global manager **420**. In some embodiments, the interface **2800** is presented on a user client **440** accessed by an administrator of the logical network. The user interface **2800** is in some embodiments a full user interface for the network management and control system, and used to define the desired state of the logical network (e.g., by issuing create, update, and delete events for logical network elements) as well as query the realized state of one or more logical network elements. The discussion below focuses on use of the interface **2800** for querying the realized state of a logical router (e.g., a Tier-0 router or a Tier-1 router). While the example in FIG. **28A** relates to a Tier-0 router for illustrative purposes, the discussion below also applies to Tier-1 routers and segments.

(205) The side panel **2805** of FIG. **28A** is a menu in some embodiments that allows the administrator to select different classes of logical network routers to query the realized state. In the example, Tier-0 Gateways have been selected in the side panel, causing the main panel **2810** to display a list of all current Tier-0 routers defined for the logical network. A single Tier-0 router is active, labeled Tier0Gateway1, and therefore highlighted by default.

(206) The main panel **2810** provides various configuration information sections for this router, including interfaces, routing, and Border Gateway Protocol (BGP) settings. In some embodiments, one or more of these sections are not configurable via the global manager as desired state, but are instead discovered state based on configuration at the physical sites.

(207) The main panel **2810** also displays in some embodiments information related to the desired state of the logical network element, such as the element's span. In the example, the Locations section **2815** shows a list of all locations that the Tier0Gateway1 router currently spans (e.g., London, Paris, and New York). Other available desired configuration information includes Domain

Name System (DNS) information, Network Address Translation (NAT) settings, and firewall configuration. Some or all of these settings may be modified by the administrator to define the desired state of the logical network.

(208) The realization status of the Tier0Gateway1 router is visually represented in some embodiments as a status indicator **2820**. In different embodiments, the status indicator may display different text, symbols, colors, or some combination of these. This status indicator **2820** represents in some embodiments a global status indicating the global realization status at all sites in the element's span.

(209) The global status indicator **2820** is in some embodiments an on-demand realization status. In some such embodiments, selecting the logical network element from the interface is the action that initiates the realization status request. In the example of FIG. **28**, only a single Tier-0 router (Tier0Gateway1) is active, so selecting the side panel item for Tier-0 routers also selects this router by default, and initiates the global realization status request. If there are multiple Tier-0 routers, then an additional selection is required in some embodiments to initiate the request, or alternatively one of the routers is selected by default and the global realization status request is issued. In other such embodiments, the global status indicator **2820** does not display any realization status even after selecting the Tier-0 router, and an additional user interaction is required to initiate the request for the global realization status. For example, in such embodiments, the global status indicator **2820** displays text or a user interface element (a button, etc.) labeled "Check status." Interaction with the indicator (e.g, the reload button to the right of the text) then initiates the global realization status request.

(210) If the router has been successfully realized at all physical sites, as illustrated in FIG. **28B**, then the global status indicator **2820** may display a colored icon (e.g., green) and/or the word "Success." If the router is not successfully realized at one or more of the physical sites in the span, then the global status indicator **2820** will display a different color (e.g., yellow or red) and/or different text such as "Pending" or "Error." In some embodiments, if the global status is not "Success" then there may be one or more alarms. The number of alarms is also displayed alongside the global status indicator in some embodiments, though in this example the number of alarms is zero since the router has been successfully realized.

(211) In some embodiments, the global status indicator **2820** is also a user interface element. Interaction with the global indicator, i.e. clicking, mouseover, or touching (if the interface is on a touchscreen) the global status indicator, creates a pop-up window in such embodiments, as in the example of FIG. **28C**. The pop-up window **2830** displays the local realization status for the logical network element at each individual physical site that is spanned by the element. In this example, each of the local managers at each spanned site each has reported "Success" to the global manager (as indicated by the status indicators **2825** next to each of the listed sites), in response to the realization status inquiry from the global manager. If the router is not successfully realized at one or more of the physical sites in the span, then the local status indicator **2825** for that site will display a different color (e.g., yellow or red) and/or different text such as "Pending" or "Error." The status may also be refreshed by selecting a refresh icon in some embodiments, either for an individual physical site or for a global refresh, which initiates a new request for realization status at the specified site or sites. If refresh is selected for a specific site, then the request includes a site identifier in some embodiments, which is used as described above with reference to operation **2607** in process **2600**.

(212) FIGS. **29A-29B** conceptually illustrate another example of a user interface **2900** for querying the realization status of logical network segments from the global manager **420**. In some embodiments, the interface **2900** is presented on a user client **440** accessed by an administrator of the logical network. The user interface **2900** is in some embodiments a full user interface for the network management and control system, and used to define the desired state of the logical network (e.g., by issuing create, update, and delete events for logical network elements) as well as

query the realized state of one or more logical network elements. The discussion below focuses on use of the interface **2900** for querying the realized state for a logical switch (e.g., an overlay segment or a VLAN segment). While the example in FIG. **29A** relates to an overlay segment for illustrative purposes, the discussion below also applies to VLAN segments.

(213) The side panel **2905** of FIG. **29A** is a menu in some embodiments that allows the administrator to select different classes of logical network routers to query the realized state. In the example, Segments have been selected in the side panel, causing the main panel **2910** to display a list of all current logical switches defined for the logical network. Multiple logical switches are active, and a logical switch labeled LS-103 is highlighted (either by default or by user interaction). Each entry provides a brief summary **2915** of the segment, indicating to which router the segment is connected, the type of segment, the associated subnet (if an overlay), and realization status. In the example, the summary indicates that logical switch LS-103 is an overlay segment associated with an IP subnet of 192.168.1.1/24, connected to router Tier0Gateway1 (described above with reference to FIG. **28**), and successfully realized.

(214) The main panel **2910** provides various configuration information sections for this segment, including segment profiles and DHCP static bindings. In some embodiments, one or more of these sections are not configurable via the global manager as desired state, but are instead discovered state based on configuration at the physical sites. Other available information includes statistics and related groups defined for the segment.

(215) The main panel also displays in some embodiments information related to the desired state of the logical network element, such as the element's span. In the example, the Locations section **2920** shows a list of all locations that the LS-103 switch currently spans (e.g., London, Paris, and New York). Other available desired configuration information **2925** includes VLAN-specific settings such as domain name and address bindings, which are not applicable for this overlay network but which are populated for VLAN segments. Some or all of these settings may be modified by the administrator to define the desired state of the logical network.

(216) The realization status of the LS-103 switch is visually represented in some embodiments as a status indicator **2930**. In different embodiments, the status indicator may display different text, symbols, colors, or some combination of these. This status indicator **2930** represents in some embodiments a global status indicating the global realization status at all sites in the element's span. For example, if the switch has been successfully realized at all physical sites, the global status indicator **2820** may display a colored icon (e.g., green) and/or the word "Success." If the switch is not successfully realized at one or more of the physical sites in the span, then the global status indicator **2820** will display a different color (e.g., yellow or red) and/or different text such as "Pending" or "Error." In some embodiments, if the global status is not "Success" then there may be one or more alarms. The number of alarms is also displayed alongside the global status indicator in some embodiments, though in this example the number of alarms is zero since the switch has been successfully realized.

(217) The global status indicator **2930** is in some embodiments an on-demand realization status. In some embodiments, selecting the logical network element from the interface is the action that initiates the realization status request. In the example of FIG. **29A**, selecting the logical switch LS-103 from the list of segments initiates the global realization status request. In some embodiments, the global status indicator does not display any realization status even after selecting the segment, and an additional user interaction is required to initiate the request for the global realization status. For example, in such embodiments, instead of displaying the global status indicator, text or a user interface element (a button, etc.) labeled "Check status" could be displayed in the same part of the interface. Interaction with this element then initiates the global realization status request.

(218) In some embodiments, the global status indicator **2930** is also a user interface element. Interaction with the global indicator, i.e. clicking, mouseover, or touching (if the interface is on a touchscreen) the global status indicator, creates a pop-up window in such embodiments, as in the

example of FIG. 29B. The pop-up window 2932 displays the local realization status for the logical network element at each individual physical site that is spanned by the element.

(219) In this example, each of the local managers at each spanned site each has reported “Success” to the global manager (as indicated by the status indicators 2935 next to each of the listed sites), in response to the realization status inquiry from the global manager. If the switch is not successfully realized at one or more of the physical sites in the span, then the local status indicator 2825 for that site will display a different color (e.g., yellow or red) and/or different text such as “Pending” or “Error.” The status may also be refreshed by selecting a refresh icon in some embodiments, either for an individual physical site or for a global refresh, which initiates a new request for realization status at the specified site or sites. If refresh is selected for a specific site, then the request includes a site identifier in some embodiments, which is used as described above with reference to operation 2607 in process 2600.

(220) In some embodiments, the global manager also receives notifications from the local managers. Unlike requests for realized state which are provided on-demand (e.g., pulled from the local managers), these notifications are automatically pushed from the local managers. The notifications pertain in some embodiments to site-specific implementation problems or situations. Notifications are provided in some embodiments to the GM as the events they describe occur, e.g. in real-time or near-real-time rather than on-demand. These notifications are also displayed in some embodiments as a report, e.g. for review by an administrator of the logical network.

(221) One type of notification in some embodiments is an alarm from a local manager that the implementation of a logical network element (e.g., a create event) has failed at the specific site. For example, the failure could be due to an IP address conflict with another network element defined locally at the local manager, of which the global manager was unaware. FIG. 30 conceptually illustrates an example of an alarm caused by a failure to create a requested logical network element at a local manager.

(222) In the first stage (upper left) of FIG. 30, the global manager 420 has a global policy tree 1700, and the local manager 430 at site B 410 has a site-specific version 2100 of the global policy tree. For clarity, only a portion of these trees are illustrated. Both trees have a Tier-0 router 3005 that was defined at the global manager 420 and successfully realized by the local manager 430. The tree at the local manager also has a Tier-1 router (T1-B) that was defined at the local manager, which references the globally-created T0 router 3005.

(223) In the second stage (upper right) of FIG. 30, a command 3015 to create a Tier-1 router 3020 (T1-G) is received at the global manager 420. Router T1-G 3010 is created under the Tier-0 router 3005 in response to the create command 3015. The span of router T1-G 3020 includes site B 410, so the global manager 420 performs a span calculation (as described above) and instructs the local manager 430 at site B to create the same router. This instruction is provided to the local manager 430 in some embodiments using process 1300, which was described above.

(224) The local manager 430 receives the instruction and attempts to fulfill the create command by creating router T1-G under the Tier-0 router. In some embodiments, the local manager attempts to create the router using process 1500, which was described above. In this example, however, the IP address range of router T1-G 3020 at least partially overlaps with the IP address range of the locally-created router T1-B 3010. Therefore, the local manager fails to create the requested router, since the requested IP range is invalid.

(225) In the third stage (bottom) of FIG. 30, the local manager 430 sends an alarm 3030 to the global manager 420 informing of the failure to create the requested router T1-G 3020 at site B 410. The alarm 3030 in some embodiments also includes relevant detailed information to diagnose the reason of the failure (e.g., the IP address range of router T1-B 3010). The local manager 430 still creates the node for the requested router in some embodiments, but does not attempt to realize the router in the control plane. The unrealized status of the router T1-G 3020 at site B 410 is visually represented by a dashed outline in FIG. 30.

(226) In some embodiments, an update command from the global manager **420** of an existing logical network element also may trigger a failure alarm. For example, the local managers **425-435** have in some embodiments emergency security policies that can override any other configuration in certain exceptional situations (e.g., a security exploit, a physical emergency or data loss, etc.). Emergency policies are not supported or recognized by the global manager in such embodiments, so any attempt to specify a change in policy that affects or contradicts the emergency policies at the local manager will also trigger an alarm in the same manner as described above with reference to FIG. **30**.

(227) Another type of notification in some embodiments is an alarm from the local manager that a local administrator of the network has attempted to override the desired state of a logical network element at that particular site. FIG. **31** conceptually illustrates an example of an alarm caused by an attempt to modify the global configuration.

(228) In the first stage (upper left left) of FIG. **31**, the global manager **420** has a global policy tree **1700**, and the local manager **430** at site B **410** has a site-specific version **2100** of the global policy tree. A Tier-0 router **3105** is defined under the global policy root **1702**, and a corresponding Tier-0 router **3110** is defined at the local manager **430** under the site-specific global policy root **2102**. For clarity, both trees are simplified versions from the counterpart examples in FIG. **17** and FIG. **21**.

(229) In the second stage (upper right) of FIG. **31**, a command **3115** to update the Tier-0 router **3110** is received at the local manager **430**. This instruction is provided to the local manager **430** in some embodiments using process **1300**, which was described above. This occurs in some embodiments by an administrator using the local manager **430** to modify the logical network rather than using the global manager **420**. The same logical network element may be modified in some embodiments by different administrators, possibly located in a different geographical regions or even countries.

(230) After receiving the update command, the local manager **430** in this example fails to implement the requested update to the Tier-0 router **3110**. One example for why the update fails is because some embodiments restrict the types of configurations that can be overridden at the local manager **430**. For example, the local manager **430** is only permitted to modify networking-related configuration in some embodiments. If the update is a security policy-related update, then the local manager is not permitted to override the global manager **420**. Another example is when the update is to extend the router's span to a new site (i.e., stretching the span of the router) prior to stretching the security policies that govern the router. Security policies must be stretched before stretching logical network elements in some embodiments, to avoid creating a security loophole after the span is extended.

(231) In the third stage (bottom) of FIG. **31**, the local manager **430** creates a shadow copy **3120** of the Tier-0 router. This shadow copy (denoted with an asterisk as Tier-0*) is initially identical to the actual T0 router **3110**, but since it was created by the local manager **430**, it may be modified according to the received update command **3115**. The original Tier-0 router **3110** from the global manager **420** is preserved as a read-only node. The local manager **430** uses the shadow copy **3120** of the router for further configuration of the network at the physical site B **410** instead of the original router **3110**.

(232) The local manager **430** sends an alarm **3125** to the global manager **420** in some embodiments, to notify it that an attempt was made to override the configuration of the Tier-0 router **3110**, and that the shadow copy **3120** is in use instead. This is necessary because without this information, the global manager **420** would assume that the original copy of the Tier-0 router **3110** is in use, and any user intent based on that assumption would potentially fail or be inconsistent.

(233) In embodiments where the attempt to modify the configuration of a logical network element succeeds (e.g., because the update is a networking-related update, not a security-related update), then the local manager in some embodiments sends a notification (not shown in FIG. **31**) to the global manager of the update. This is necessary to inform the global manager **420** that the realized

state of the logical network element at this physical site will not match the realized state of the element at other sites, due to the site-specific update.

(234) A third type of notification in some embodiments is a message from the local manager that the deletion of a logical network element has failed at the site. FIGS. 32A-32B conceptually illustrate an example of an alarm caused by a failure to delete a logical network element.

(235) In the first stage of the example (upper portion of FIG. 32A), the global manager 420 has a global policy tree 1700, and local managers 425 and 430 at sites A 405 and B 410 respectively have a site-specific versions 2100 of the global policy tree. At the global manager 420, a Tier-0 router 3205 is defined under the global policy root 1702, and corresponding Tier-0 routers 3210 and 3215 are defined at the local managers 425 and 430 under the corresponding site-specific global policy roots 3220 and 3225. For clarity, all three trees are simplified versions from the counterpart examples in FIG. 17 and FIG. 21.

(236) In the example, at site A 405 the local manager 425 also has a local policy tree with local root 3230 which was defined locally (e.g., by a local administrator, who may be different from a global administrator defining the global policy tree at the global manager 420). Beneath the local root 3230, a local Tier 1 router 3235 is also defined. This local Tier-1 router has a reference to the globally-defined Tier-0 router 3210 from the global manager (depicted as a dashed arrow from Tier-1 router 3235 to Tier-0 router 3210). The global manager 420 is completely unaware of the existence of Tier-1 router 3235 in some embodiments.

(237) At site B 410, the local manager 430 also has a local policy tree with local root 3240 which was defined locally (e.g., by a local administrator, who may be different from a global administrator defining the global policy tree at the global manager 420). Beneath the local root 3240, a local Tier 1 router 3245 is also defined. This local Tier-1 router 3245 is a completely different router than the local Tier-1 router 3235 defined beneath the local root 3230 at the local manager 425 for site A 405. The local Tier-1 router 3245 does not have any reference to the globally-defined Tier-0 router 3210 from the global manager.

(238) In the second stage of the example (bottom of FIG. 32A), a command 3250 to delete the Tier-0 router 3205 is received at the global manager 420. However, the Tier-0 router 3205 is not immediately deleted at the global manager in response to the command 3250, in some embodiments. The global manager 420 performs a span calculation to identify sites A and B as the span of the router 3205, and sends an instruction to the local managers 425 and 430, to delete the realized Tier-0 routers 3210 and 3215. This instruction is provided to the local manager 430 in some embodiments using process 1300, which was described above.

(239) In the third stage of the example (top of FIG. 32B), the local manager 425 at site A 405 receives the instruction, but fails to delete the Tier-0 router 3210, since it is being referenced by the local Tier-1 router 3235. The local manager 430 at site B 410 also receives the instruction and successfully deletes the Tier-0 router 3215, since there are no local references to the router at this site.

(240) In the fourth stage of the example (bottom of FIG. 32B), the local manager 425 sends an alarm 3255 to the global manager 420 informing of the failure to delete the requested Tier-0 router at site A 405. The alarm 3255 in some embodiments also includes detailed information to diagnose the reason of the failure. The Tier-0 router 3260 is not deleted, but is marked as pending deletion by the global manager in some embodiments, as illustrated by a dashed outline. In some embodiments an additional notification is generated for the local administrator of the local manager 425, so that they can resolve the issue (for example, by changing the reference of the T-1 router 3235 to a different T0 router, such as one created at the local manager 425).

(241) The local manager 430 at site B 410 sends a notification 3265 to the global manager to inform that the Tier-0 router 3215 was successfully deleted from the local root 3225. The global manager 420 receives the notification 3265, but since there is an alarm 3255 instead of a corresponding success notification from local manager 425, in some embodiments the Tier-0 router

3205 is not deleted from the global root **1702**.

(242) An additional example of a type of notification in some embodiments is a message from the local manager that a logical network endpoint (e.g., a virtual machine) has been migrated from one site to another site. Such endpoints are typically attached to a logical port of a logical network element (e.g., a logical switch) that spans both sites. When the endpoint is migrated to a different site, it retains its association with that same logical port in some embodiments, and seamlessly inherits the same security policies that are associated with that port, even at the new site.

(243) In some embodiments, a machine may be migrated from one physical site to another physical site, while preserving the state of network policies that apply to the machine. For example, for a logical segment (logical switch) that spans at least two physical sites, a machine attached to a logical port of the segment at one of the spanned sites is migrated in some embodiments to a new logical port of the same segment at another one of the spanned sites. Migrating a machine in some embodiments comprises creating a new machine at a new host computer (at the new site), transferring configuration data to the new machine from the original machine, and removing the original machine from its host computer. In some embodiments, the configuration data comprises physical state data and logical network policies.

(244) The state of the machine (e.g., runtime data) is transferred in some embodiments from a computing device executing the original machine at the source site (i.e., the source host) to a computing device executing the new machine at the destination site (i.e., the destination host). In some embodiments, the state includes a local port identifier assigned by a local manager (e.g., local manager **425**) at the source physical site, and a global port identifier assigned by the global manager **420**. The global port identifier is associated with logical network policies defined through a user client **440** (e.g., by an administrator of the network) in some embodiments as part of the desired configuration of the logical network. Upon receiving the transferred state data, the destination host overwrites at least a portion of the state of the new machine with the received state of the old machine.

(245) In some embodiments, the destination host extracts the global port identifier from the received state data and provides it to the destination site's local manager (e.g., local manager **430**). The destination local manager **430** uses the global port identifier in some embodiments to retrieve logical network policies associated with the global port identifier from the global manager **420**, and apply these policies to the new machine. In some embodiments, the local manager **430** generates configuration data from the policies and provides the configuration data to the control plane of the logical network at the destination site (e.g., a set of controllers **1420-1425**) which then distributes the configuration data to the destination host.

(246) FIG. **33** illustrates an example of some embodiments for migrating a virtual machine between sites within a region while preserving the machine's group membership and associated policies. The virtual machine (VM) **3305** in this example is a member of group D **1825** (e.g., machines running OS2 in the Europe domain) as described above with reference to FIG. **18**. In this example, a simplified version of the global policy tree **3310** at the global manager **420** is shown, along with the corresponding (simplified) global policy trees **3315-3320** at the source and destination sites' local managers **425-430** (e.g., Paris and London, respectively). The policy tree has a Tier 1 router **3325** that spans both sites, and a network segment **3330** that references (dashed arrow) the T1 router **3325** as a parent and therefore inherits the same span. As a result, the router **3325** and the segment **3330** are defined in the policy trees at both sites. The segment **3330** also has a logical port **3327**, which is assigned a global port identifier LP-1 at the global manager **420**.

(247) The initial stage of FIG. **33** (top half) illustrates the policy trees prior to the migration. Here, VM C **3305** has been created and attached to a logical port **3336** of segment **3330** at the first physical site A (e.g., Paris **405**). VM C **3305** is a member of group D **1805**, and executes on one of the host computers **1805** at site A. The port **3336** is assigned a local port identifier LP-A (e.g., by the local manager at site A) that in some embodiments is specific to the source site. A set of policies

(not shown) are configured to apply to group D **1825**, and are therefore applied to VM C **3305** at site A **405**, since the domain **1770** under which group D is defined spans site A. A migration is then initiated (dotted line) to transfer VM C **3305** from a source server **1805** at site A (**405**) to a destination server **1810** at site B (**410**).

(248) The second stage of FIG. **33** (bottom half) illustrates the policy trees after the migration. Here, VM C **3305** has been successfully transferred to the second physical site B (e.g., London **410**). VM C **3305** is now attached to a logical port **3337** of segment **3330** at the second physical site B (e.g., London **410**). VM C **3305** remains a member of group D **1805**, and now executes on one of the host computers **1810** at site B. The port **3337** is assigned a local port identifier LP-B (e.g., by the local manager at site B) that in some embodiments is specific to the destination site. The set of policies remain applied to VM C **3305** at site B **410**, since the domain **1770** under which group D is defined spans site B.

(249) In this example, the global manager is not aware of the local port identifiers LP-A or LP-B. In some embodiments, each of the local managers **425-430** maintains a mapping of the global port identifier and the local port identifier for the logical port at the respective site. This mapping allows the local managers to retrieve from the global manager **420** the relevant policy information applicable to the machine, so that these policies are seamlessly applied before and after migration.

(250) In the example of FIG. **33**, the machine was transferred between sites in the same region, retained its group membership, and policies associated with that group membership were applied automatically. Some embodiments also permit migration of a machine from a site in one region to a site in another region. However, in some such embodiments, the same policies would not necessarily apply to the machine (e.g., if the policies applicable to the machine at the source site do not span the second site). In such cases, in some embodiments new security policies would need to be applied manually (e.g., by the administrator of the logical network). In addition, if the segment does not span the destination site, then in some embodiments the machine would lose connectivity to the rest of the logical network after migration.

(251) FIGS. **34A-C** conceptually illustrate the exchange of information during VM creation and migration in some embodiments between the global manager, the local managers, and source and destination hosts. FIG. **35** conceptually illustrates, with reference to FIG. **34**, a process **3500** performed in some embodiments by a source host **3415** at the source site during creation and migration of a virtual machine (VM). In embodiments exemplified by this process, the VM is migrated from the source site by powering on a new VM at the destination site, and transferring the VM state from the source host to the destination host.

(252) The process **3500** begins at **3505** by receiving an instruction from a compute manager **3405** to register a new VM, for attachment to a logical segment **3330** that spans the source site. The register command is illustrated in FIG. **34A** as instruction **3422**. The compute manager **3405** in some embodiments directs the allocation of computing resources at one or more of the physical sites spanned by the logical network. In some embodiments, an administrator of the logical network determines (e.g., through a user interface **440**) where the new VM should be registered.

(253) At **3510**, the process **3500** requests a port attachment for the new VM, from the local manager **425** at the source site. The port attachment request is illustrated in FIG. **34A** as instruction **3424**. In response to the request, the local manager **425** creates a logical port **3336** for the logical segment **3330**, and returns a local port identifier (e.g., LP-A) specific to the source site, as illustrated in FIG. **34A** by instruction **3426**. The local port identifier is a universal unique identifier (UUID) in some embodiments.

(254) At **3515**, the process **3500** powers on the requested VM, as indicated by event **3427**. The process **3500** attaches the new VM to the port **3336** of the logical segment **3330**, using the local port identifier (e.g., LP-A) received from the local manager **425**. In some embodiments, the process **3500** sends a success message to the local manager **425** after attaching the port **3336**. The success message is illustrated in FIG. **34A** as instruction **3428**.

(255) Upon receiving the success message, in some embodiments the local manager **425** provides the local port identifier (as illustrated in FIG. **34A** by instruction **3432**) to the global manager **420** (e.g. discovered state, as described above). The global manager updates the segment **3330** to include a logical port **3327**, and allocates a global port identifier (e.g., LP-1) to the logical port. In some embodiments, the global manager **420** uses the same local port identifier assigned by the source local manager **425**. In other embodiments, the global manager generates a new port identifier, e.g. a path relative to the global root **1702** that indicates the location of the parent segment in the hierarchical policy tree. The global manager **3425** then provides the global port identifier back to the local manager **425**, as illustrated in FIG. **34A** by instruction **3434**.

(256) At **3520**, the process **3500** receives the global port identifier from the local manager **425**, as illustrated in FIG. **34A** by instruction **3436**. In some embodiments, the local manager **425** provides the global identifier to the central control plane of the logical network at the source site (e.g., a set of controllers **1420-1425**). The controllers then distribute the global identifier to the process **3500**, via an agent in some embodiments that executes on the same host (e.g., the source host **3415** that performs the process **3500**) alongside the VM.

(257) In some embodiments, a set of tags are also specified at the global manager **420** for attaching to the newly attached VM. These tags are used to define group membership for the VM in some embodiments. Security rules may also be defined dynamically using these tags. In some embodiments, these dynamic references to tags may include regular expressions. Thus, for example, a set of tags `dev_paris` could be defined to indicate a VM is in development mode, and security rules defined just for the Paris site by referencing the full tag, or security rules could be defined just using a portion of the tag (e.g., “`dev_*`”) to apply to all VMs tagged as in development mode regardless of location. The tags are directly applied to the VM in some embodiments, and applied by the source host **3415** internally on the logical port.

(258) In some embodiments, the global manager **420** receives configuration data (e.g., desired state, such as security policies) to apply to the logical port. The desired state is received in some embodiments through a user interface **440** (e.g., from an administrator of the logical network). The desired state is then provided from the global manager **420** to the local manager **425**, as described above with reference to FIG. **13**. The receipt of the desired state by the local manager **425** is illustrated in FIG. **34A** as instruction **3438**.

(259) At **3525**, the process **3500** receives and applies the desired configuration (e.g., security policies, of which some may reference the tags) to the logical port **3336**, as illustrated in FIG. **34A** by instruction **3442**. In other words, the process **3500** realizes the desired state. In some embodiments, the local manager **425** generates configuration data from the received desired state, and provides the configuration data to the central control plane of the logical network at the source site (e.g., a set of controllers **1420-1425**). The controllers then distribute the configuration data to the process **3500**, via an agent in some embodiments that executes on the same host (e.g., the source host **3415** that performs the process **3500**) alongside the VM.

(260) At **3530**, the process **3500** receives an instruction from the compute manager **3405** to prepare the VM for migration. The migration command is illustrated in FIG. **34A** as instruction **3444**. The instruction includes in some embodiments an address of the destination host **3410** to where the VM will migrate. In some embodiments, an administrator of the logical network determines (e.g., through a user interface **440**) where the VM will migrate.

(261) At **3535**, the process **3500** saves the realized state data for the VM, as indicated by event **3445**. In some embodiments, the saved state data includes the global port identifier, which was received as described above at **3520**. The process **3500** sends the saved data to the destination host **3410**, as described below with reference to process **3600** of FIG. **36**, and as illustrated in FIG. **34B** as instruction **3458**.

(262) At **3540**, the process **3500** receives an instruction from the compute manager **3405** to de-register the VM. The de-registration command is illustrated in FIG. **34C** as instruction **3446**. In

some embodiments, the compute manager **3405** sends this instruction based on one or more notifications (not shown) that the state data has been successfully saved (e.g., from the local manager **425**) and/or that the VM at the destination site has been successfully updated with the saved state data, as described below with reference to process **3600** of FIG. **36**.

(263) At **3545**, the process **3500** requests the local manager **425** to detach the logical port **3336**, as illustrated in FIG. **34C** by instruction **3448**. After successfully detaching the port, the process **3500** powers off the VM, as indicated by event **3450**. The process **3500** then ends.

(264) FIG. **36** conceptually illustrates, with reference to FIG. **34**, a corresponding process **3600** performed in some embodiments by a destination host **3410** at the destination site during migration of the VM from the source host **3415**. In embodiments exemplified by this process, the VM is migrated from the source site by powering on a new VM at the destination site, and transferring the VM state from the source host to the destination host.

(265) The process **3600** begins at **3605** by receiving an instruction from the compute manager **3405** to register a new VM, for attachment to the logical segment **3330** that spans the source site and the destination site. The register command is illustrated in FIG. **34B** as instruction **3452**. As noted above, the compute manager **3405** in some embodiments directs the allocation of computing resources at one or more of the physical sites spanned by the logical network. In some embodiments, an administrator of the logical network determines (e.g., through a user interface **440**) to which site the VM should be migrated.

(266) At **3610**, the process **3600** requests a port attachment for the new VM, from the local manager **430** at the destination site. The port attachment request is illustrated in FIG. **34B** as instruction **3454**. In response to the request, the local manager **430** creates a logical port **3337** for the logical segment **3330**, and returns a local port identifier (e.g., LP-B) specific to the source site, as illustrated in FIG. **34B** by instruction **3456**. In some embodiments, the local manager **430** also includes flag for the created port **3337**, to indicate that the port has a pending state.

(267) At **3615**, the process **3600** powers on the requested VM, as indicated in FIG. **34B** by event **3457**. The process **3600** attaches the new VM to the port **3337** of the logical segment **3330**, using the local port identifier (e.g., LP-B) received from the local manager **430**.

(268) At **3620**, the process **3600** receives saved state data from the source host **3415**. The saved state data includes the global port identifier (e.g., LP-1) in some embodiments, which was saved by the original VM on the source host **3415**, as described above with reference to process **3500** of FIG. **35**. The receipt of the state data and global identifier is illustrated in FIG. **34B** as instruction **3458**.

(269) In some embodiments, the transfer **3458** of state data from the source host **3415** is a lengthy process, including copying disks, memory, and other data. This process takes several hours in some embodiments, and the source VM is still active on the source host **3415** during this time. Packets received by the source host are also included in the state transfer, so even though there are no policies applied to the destination host during the state transfer, it does not matter, since the traffic is being correctly processed by the source VM. The destination VM does not receive any direct data traffic during the state transfer. Therefore, in some embodiments, the pending flag is used to ensure that the destination host does not send a success message to the local manager until the state transfer is complete. Otherwise, the local manager would report the success to the global manager and traffic would begin to arrive prematurely at the destination VM. The pending state is tracked in some embodiments by both the destination host **3410** and the destination local manager **430**.

(270) At **3625**, the process **3600** overwrites the configuration of the logical port **3337** with the received state data, as indicated in FIG. **34B** by event **3459**. Specifically, the process **3600** restores the saved data from the original VM to the new VM. The state data also includes in some embodiments the local port identifier (e.g., LP-A) assigned to the original VM from the source local manager **430**, as described above with reference to FIG. **35**. In such embodiments, the process does not overwrite the local port identifier (e.g., LP-B) received from the local manager **430** with

the local port identifier (e.g., LP-A) in the saved data. In some embodiments, the state data also includes metadata associated with the original VM, some of which is also overwritten onto the new VM. The state data also includes any tags that were assigned to the source VM.

(271) In some embodiments, when the destination local manager **430** identifies that the full replication is complete (memory, disks, etc), it performs a fast suspend and resume (FSR) operation at the destination VM. This is coordinated with the source local manager **425** in some embodiments (e.g., via the compute manager **3405**), to also perform the same FSR operation at the source VM. The FSR operation flips active status of the source VM off and enables active status for the destination VM, so that the destination VM may begin receiving traffic. The pending state of the destination port is also removed in some embodiments at this stage.

(272) At **3630**, the process **3600** sends a success message to the local manager **430** after attaching the port. In some embodiments, the process **3600** also reports the global port identifier (e.g., LP-1), and also any associated tags, to the local manager **430**. The provision of the success message, global port identifier, and tags to the local manager **430** is illustrated in FIG. **34B** as instruction **3462**.

(273) Upon receiving the global port identifier, in some embodiments the local manager **430** updates the segment **3330** to associate the global port identifier (e.g., LP-1) with the local port identifier (e.g., LP-B) for port **3337**. The local manager **430** then uses the global port identifier and the tags to request the desired state for the port **3337**, from the global manager **420**, as illustrated in FIG. **34B** by instruction **3464** to the global manager **420**. The desired state is then provided from the global manager **420** to the local manager **430**, as described above with reference to FIG. **13**. The receipt of the desired state by the local manager is illustrated in FIG. **34B** as instruction **3466**.

(274) At **3635**, the process **3600** receives and applies the desired state to the logical port **3337**, as illustrated in FIG. **34B** by instruction **3468**. In other words, at **3635** the process **3600** realizes the desired state. In some embodiments, the local manager **430** generates configuration data from the received desired state, and provides the configuration data to the central control plane of the logical network at the source site (e.g., a set of controllers **1420-1425**). The controllers then distribute the configuration data to the process **3600**. The process **3600** then ends.

(275) Many of the above-described features and applications are implemented as software processes that are specified as a set of instructions recorded on a computer readable storage medium (also referred to as computer readable medium). When these instructions are executed by one or more processing unit(s) (e.g., one or more processors, cores of processors, or other processing units), they cause the processing unit(s) to perform the actions indicated in the instructions. Examples of computer readable media include, but are not limited to, CD-ROMs, flash drives, RAM chips, hard drives, EPROMs, etc. The computer readable media does not include carrier waves and electronic signals passing wirelessly or over wired connections.

(276) In this specification, the term “software” is meant to include firmware residing in read-only memory or applications stored in magnetic storage, which can be read into memory for processing by a processor. Also, in some embodiments, multiple software inventions can be implemented as sub-parts of a larger program while remaining distinct software inventions. In some embodiments, multiple software inventions can also be implemented as separate programs. Finally, any combination of separate programs that together implement a software invention described here is within the scope of the invention. In some embodiments, the software programs, when installed to operate on one or more electronic systems, define one or more specific machine implementations that execute and perform the operations of the software programs.

(277) FIG. **37** conceptually illustrates an electronic system **3700** with which some embodiments of the invention are implemented. The electronic system **3700** may be a computer (e.g., a desktop computer, personal computer, tablet computer, server computer, mainframe, a blade computer etc.), phone, PDA, or any other sort of electronic device. Such an electronic system includes various types of computer readable media and interfaces for various other types of computer readable

media. Electronic system **3700** includes a bus **3705**, processing unit(s) **3710**, a system memory **3725**, a read-only memory **3730**, a permanent storage device **3735**, input devices **3740**, and output devices **3745**.

(278) The bus **3705** collectively represents all system, peripheral, and chipset buses that communicatively connect the numerous internal devices of the electronic system **3700**. For instance, the bus **3705** communicatively connects the processing unit(s) **3710** with the read-only memory **3730**, the system memory **3725**, and the permanent storage device **3735**.

(279) From these various memory units, the processing unit(s) **3710** retrieve instructions to execute and data to process in order to execute the processes of the invention. The processing unit(s) may be a single processor or a multi-core processor in different embodiments.

(280) The read-only-memory (ROM) **3730** stores static data and instructions that are needed by the processing unit(s) **3710** and other modules of the electronic system. The permanent storage device **3735**, on the other hand, is a read-and-write memory device. This device is a non-volatile memory unit that stores instructions and data even when the electronic system **3700** is off. Some embodiments of the invention use a mass-storage device (such as a magnetic or optical disk and its corresponding disk drive) as the permanent storage device **3735**.

(281) Other embodiments use a removable storage device (such as a floppy disk, flash drive, etc.) as the permanent storage device. Like the permanent storage device **3735**, the system memory **3725** is a read-and-write memory device. However, unlike storage device **3735**, the system memory is a volatile read-and-write memory, such a random-access memory. The system memory stores some of the instructions and data that the processor needs at runtime. In some embodiments, the invention's processes are stored in the system memory **3725**, the permanent storage device **3735**, and/or the read-only memory **3730**. From these various memory units, the processing unit(s) **3710** retrieve instructions to execute and data to process in order to execute the processes of some embodiments.

(282) The bus **3705** also connects to the input and output devices **3740** and **3745**. The input devices enable the user to communicate information and select commands to the electronic system. The input devices **3740** include alphanumeric keyboards and pointing devices (also called “cursor control devices”). The output devices **3745** display images generated by the electronic system. The output devices include printers and display devices, such as cathode ray tubes (CRT) or liquid crystal displays (LCD). Some embodiments include devices such as a touchscreen that function as both input and output devices.

(283) Finally, as shown in FIG. **37**, bus **3705** also couples electronic system **3700** to a network **3765** through a network adapter (not shown). In this manner, the computer can be a part of a network of computers (such as a local area network (“LAN”), a wide area network (“WAN”), or an Intranet, or a network of networks, such as the Internet. Any or all components of electronic system **3700** may be used in conjunction with the invention.

(284) Some embodiments include electronic components, such as microprocessors, storage and memory that store computer program instructions in a machine-readable or computer-readable medium (alternatively referred to as computer-readable storage media, machine-readable media, or machine-readable storage media). Some examples of such computer-readable media include RAM, ROM, read-only compact discs (CD-ROM), recordable compact discs (CD-R), rewritable compact discs (CD-RW), read-only digital versatile discs (e.g., DVD-ROM, dual-layer DVD-ROM), a variety of recordable/rewritable DVDs (e.g., DVD-RAM, DVD-RW, DVD+RW, etc.), flash memory (e.g., SD cards, mini-SD cards, micro-SD cards, etc.), magnetic and/or solid state hard drives, read-only and recordable Blu-Ray® discs, ultra-density optical discs, any other optical or magnetic media, and floppy disks. The computer-readable media may store a computer program that is executable by at least one processing unit and includes sets of instructions for performing various operations. Examples of computer programs or computer code include machine code, such as is produced by a compiler, and files including higher-level code that are executed by a computer,

an electronic component, or a microprocessor using an interpreter.

(285) While the above discussion primarily refers to microprocessor or multi-core processors that execute software, some embodiments are performed by one or more integrated circuits, such as application specific integrated circuits (ASICs) or field programmable gate arrays (FPGAs). In some embodiments, such integrated circuits execute instructions that are stored on the circuit itself.

(286) As used in this specification, the terms “computer”, “server”, “processor”, and “memory” all refer to electronic or other technological devices. These terms exclude people or groups of people. For the purposes of the specification, the terms display or displaying means displaying on an electronic device. As used in this specification, the terms “computer readable medium,” “computer readable media,” and “machine readable medium” are entirely restricted to tangible, physical objects that store information in a form that is readable by a computer. These terms exclude any wireless signals, wired download signals, and any other ephemeral signals.

(287) This specification refers throughout to computational and network environments that include virtual machines (VMs). However, virtual machines are merely one example of data compute nodes (DCNs) or data compute end nodes, also referred to as addressable nodes. DCNs may include non-virtualized physical hosts, virtual machines, containers that run on top of a host operating system without the need for a hypervisor or separate operating system, and hypervisor kernel network interface modules.

(288) VMs, in some embodiments, operate with their own guest operating systems on a host using resources of the host virtualized by virtualization software (e.g., a hypervisor, virtual machine monitor, etc.). The tenant (i.e., the owner of the VM) can choose which applications to operate on top of the guest operating system. Some containers, on the other hand, are constructs that run on top of a host operating system without the need for a hypervisor or separate guest operating system. In some embodiments, the host operating system uses name spaces to isolate the containers from each other and therefore provides operating-system level segregation of the different groups of applications that operate within different containers. This segregation is akin to the VM segregation that is offered in hypervisor-virtualized environments that virtualize system hardware, and thus can be viewed as a form of virtualization that isolates different groups of applications that operate in different containers. Such containers are more lightweight than VMs.

(289) Hypervisor kernel network interface modules, in some embodiments, is a non-VM DCN that includes a network stack with a hypervisor kernel network interface and receive/transmit threads. One example of a hypervisor kernel network interface module is the vmknix module that is part of the ESXi™ hypervisor of VMware, Inc.

(290) It should be understood that while the specification refers to VMs, the examples given could be any type of DCNs, including physical hosts, VMs, non-VM containers, and hypervisor kernel network interface modules. In fact, the example networks could include combinations of different types of DCNs in some embodiments.

(291) While the invention has been described with reference to numerous specific details, one of ordinary skill in the art will recognize that the invention can be embodied in other specific forms without departing from the spirit of the invention. In addition, a number of the figures (including FIGS. 10, 12, 13, 15, 16, 19, 20, 24, 26, 35, and 36) conceptually illustrate processes. The specific operations of these processes may not be performed in the exact order shown and described. The specific operations may not be performed in one continuous series of operations, and different specific operations may be performed in different embodiments. Furthermore, the process could be implemented using several sub-processes, or as part of a larger macro process. Thus, one of ordinary skill in the art would understand that the invention is not to be limited by the foregoing illustrative details, but rather is to be defined by the appended claims.

Claims

1. A method for defining a logical network that spans a plurality of sites and is managed at each site by a local manager, the method comprising: at a global manager that manages the logical network: storing a definition of the logical network that spans the plurality of sites, the definition specifying (i) a plurality of logical network elements, (ii) relationships between the elements, and (iii) for each of at least a subset of the elements, a span attribute indicating a set of sites spanned by the element; based on the span attributes and relationships between the elements, identifying for each respective site a portion of the logical network definition that is relevant for the respective site; and providing, to the local manager of each respective site, the portion of the logical network definition that is identified for the respective site, the local manager being configured to apply a set of priority rules to determine an override of a global configuration by a local configuration update event; wherein: the global manager comprises a primary global manager and a secondary global manager, the primary global manager being asynchronously synchronized with the secondary global manager as a standby for failover scenarios, the secondary global manager being configured at a different physical site than the primary global manager and maintains an independent database; a particular local manager at a particular site creates and stores a local definition of the logical network based on the logical network definition portion received from the global manager; the portion received by the particular local manager from the global manager is a first logical network definition; and the particular local manager directly receives a second logical network definition from a user and updates the local definition of the logical network based on the second logical network definition; the particular local manager stores the logical network definition portion received from the global manager as a read-only definition that is not modified when updating the local definition of the logical network.
2. The method of claim 1, wherein: a particular local manager at a particular site uses the portion of the logical network definition identified for the particular site to generate and provide a set of configuration data to a cluster of controllers at the particular site; and the cluster of controllers (i) receives the configuration data from the particular local manager, (ii) identifies a set of computing devices at the particular site to implement the logical network elements that span the particular site, and (ii) distributes the configuration data to the identified set of computing devices.
3. The method of claim 2, wherein: the configuration data received by a particular computing device is received by a local controller executing on the particular computing device; and the local controller uses the received configuration data to configure at least one forwarding element also executing on the particular computing device to implement at least one of the logical network elements that spans the particular site.
4. The method of claim 1, wherein the definition of the logical network further specifies a plurality of logical network policies, including forwarding policies, security policies, and service policies.
5. The method of claim 1 further comprising determining a span attribute for a first logical network element based on (i) a defined span attribute of a second logical network element and (ii) a relationship between the first and second logical network elements.
6. The method of claim 1, wherein providing a particular identified portion of the logical network definition to a particular local manager comprises writing the particular identified portion of the logical network definition to a file and transmitting the file to the particular local manager over a network.
7. The method of claim 1, wherein the span attribute for a particular logical network element in the logical network definition comprises a set of site identifiers that are each associated with one of the plurality of sites.
8. A non-transitory machine-readable medium storing a global manager program which when executed by at least one processing unit defines a logical network that spans a plurality of sites and is managed at each site by a local manager, the program comprising sets of instructions for: storing a definition of the logical network that spans the plurality of sites, the definition specifying (i) a

plurality of logical network elements, (ii) relationships between the elements, and (iii) for each of at least a subset of the elements, a span attribute indicating a set of sites spanned by the element; based on the span attributes and relationships between the elements, identifying for each respective site a portion of the logical network definition that is relevant for the respective site; and providing, to the local manager of each respective site, the portion of the logical network definition that is identified for the respective site, the local manager being configured to apply a set of priority rules to determine an override of a global configuration by a local configuration update event; wherein: the global manager comprises a primary global manager and a secondary global manager, the primary global manager being asynchronously synchronized with the secondary global manager as a standby for failover scenarios, the secondary global manager being configured at a different physical site than the primary global manager and maintains an independent database; a particular local manager at a particular site creates and stores a local definition of the logical network based on the logical network definition portion received from the global manager; the portion received by the particular local manager from the global manager is a first logical network definition; and the particular local manager directly receives a second logical network definition from a user and updates the local definition of the logical network based on the second logical network definition; the particular local manager stores the logical network definition portion received from the global manager as a read-only definition that is not modified when updating the local definition of the logical network.

9. The non-transitory machine-readable medium of claim 8, wherein: a particular local manager at a particular site uses the portion of the logical network definition identified for the particular site to generate and provide a set of configuration data to a cluster of controllers at the particular site; and the cluster of controllers (i) receives the configuration data from the particular local manager, (ii) identifies a set of computing devices at the particular site to implement the logical network elements that span the particular site, and (i) distributes the configuration data to the identified set of computing devices.

10. The non-transitory machine-readable medium of claim 9, wherein: the configuration data received by a particular computing device is received by a local controller executing on the particular computing device; and the local controller uses the received configuration data to configure at least one forwarding element also executing on the particular computing device to implement at least one of the logical network elements that spans the particular site.

11. The non-transitory machine-readable medium of claim 8, wherein the definition of the logical network further specifies a plurality of logical network policies, including forwarding policies, security policies, and service policies.

12. The non-transitory machine-readable medium of claim 8, wherein the global manager program further comprises a set of instructions for determining a span attribute for a first logical network element based on (i) a defined span attribute of a second logical network element and (ii) a relationship between the first and second logical network elements.

13. The non-transitory machine-readable medium of claim 8, wherein the set of instructions for providing a particular identified portion of the logical network definition to a particular local manager comprises a set of instructions for writing the particular identified portion of the logical network definition to a file and transmitting the file to the particular local manager over a network.

14. The non-transitory machine-readable medium of claim 8, wherein the span attribute for a particular logical network element in the logical network definition comprises a set of site identifiers that are each associated with one of the plurality of sites.
