

# US Patent & Trademark Office

## Patent Public Search | Text View

---

|                      |                        |
|----------------------|------------------------|
| United States Patent | 12399913               |
| Kind Code            | B1                     |
| Date of Patent       | August 26, 2025        |
| Inventor(s)          | Toulme; Antoine et al. |

---

### Versioned connectors for data ingestion from distributed ledger systems

---

#### Abstract

Systems and methods for ingesting data from distributed ledger systems. To ingest data from the distributed ledger systems, all or a portion of the distributed ledger systems can be assigned a connector. One or more connectors can be initialized for a particular distributed ledger system. For example, the data intake and query system can include a first connector that is initialized for a first distributed ledger system and a second connector that is initialized for a second distributed ledger system. The connectors can be initiated with processing schema that can indicate how the connectors ingest, enrich, filter, etc. data from the distributed ledger systems. All or a portion of the connectors can be initiated with different processing schemas. The connectors can obtain data from nodes of the distributed ledger systems according to the processing schemas and route the data to the data intake and query system.

---

**Inventors:** Toulme; Antoine (San Jose, CA), Marsollier; Melanie (San Francisco, CA), McKerverey; Nathaniel Gerard (Tallahassee, FL), Moore; Ryan Russell (St. Augustine, FL)

**Applicant:** Splunk LLC (San Francisco, CA)

**Family ID:** 1000006561938

**Assignee:** Cisco Technology, Inc. (San Jose, CA)

**Appl. No.:** 17/816151

**Filed:** July 29, 2022

#### Related U.S. Application Data

us-provisional-application US 63366134 20220609

---

#### Publication Classification

**Int. Cl.: G06F7/00** (20060101); **G06F16/21** (20190101); **G06F16/28** (20190101); **H04L9/00** (20220101)

**U.S. Cl.:**

**CPC G06F16/287** (20190101); **G06F16/211** (20190101); **H04L9/50** (20220501);

## Field of Classification Search

**CPC:** G06F (16/287); G06F (16/211); G06F (16/2379); G06F (16/245); G06F (16/248); H04L (9/50); H04L (9/3247); H04L (9/0891)

**USPC:** 707/703; 707/722; 707/736; 707/737; 707/748; 707/688; 707/769; 707/770; 707/802; 705/7.11; 705/7.12; 705/7.13; 705/7.17; 705/7.22; 705/7.29; 705/7.35; 705/28; 705/37; 700/29; 700/43; 700/117; 700/118; 700/119; 700/248

---

## References Cited

### U.S. PATENT DOCUMENTS

| Patent No.   | Issued Date | Patentee Name    | U.S. Cl. | CPC          |
|--------------|-------------|------------------|----------|--------------|
| 7937344      | 12/2010     | Baum et al.      | N/A      | N/A          |
| 8112425      | 12/2011     | Baum et al.      | N/A      | N/A          |
| 8751529      | 12/2013     | Zhang et al.     | N/A      | N/A          |
| 8788525      | 12/2013     | Neels et al.     | N/A      | N/A          |
| 8965902      | 12/2014     | Navas            | N/A      | N/A          |
| 9215240      | 12/2014     | Merza et al.     | N/A      | N/A          |
| 9286413      | 12/2015     | Coates et al.    | N/A      | N/A          |
| 9825931      | 12/2016     | Johnsrud et al.  | N/A      | N/A          |
| 10127258     | 12/2017     | Lamas et al.     | N/A      | N/A          |
| 10402792     | 12/2018     | Lin et al.       | N/A      | N/A          |
| 10419469     | 12/2018     | Singh et al.     | N/A      | N/A          |
| 10438133     | 12/2018     | Vachhani et al.  | N/A      | N/A          |
| 11062042     | 12/2020     | McKervery        | N/A      | H04L 63/123  |
| 11269859     | 12/2021     | Luedtke          | N/A      | G06F 16/245  |
| 11356243     | 12/2021     | Umemoto          | N/A      | H04L 63/0861 |
| 11507562     | 12/2021     | Luedtke et al.   | N/A      | N/A          |
| 11921873     | 12/2023     | McKervery et al. | N/A      | N/A          |
| 2003/0182310 | 12/2002     | Charnock et al.  | N/A      | N/A          |
| 2009/0313638 | 12/2008     | Shetty           | N/A      | N/A          |
| 2015/0019537 | 12/2014     | Neels et al.     | N/A      | N/A          |
| 2015/0356175 | 12/2014     | Flynn et al.     | N/A      | N/A          |
| 2018/0018590 | 12/2017     | Szeto            | N/A      | N/A          |
| 2018/0082024 | 12/2017     | Curbera et al.   | N/A      | N/A          |
| 2018/0082043 | 12/2017     | Witchey et al.   | N/A      | N/A          |
| 2018/0113578 | 12/2017     | Yoon et al.      | N/A      | N/A          |
| 2018/0189781 | 12/2017     | McCann et al.    | N/A      | N/A          |
| 2018/0268401 | 12/2017     | Ortiz            | N/A      | N/A          |
| 2019/0079950 | 12/2018     | Ramabaja         | N/A      | N/A          |
| 2019/0098106 | 12/2018     | Mungel et al.    | N/A      | N/A          |
| 2019/0109713 | 12/2018     | Clark et al.     | N/A      | N/A          |

|              |         |                    |     |             |
|--------------|---------|--------------------|-----|-------------|
| 2019/0121566 | 12/2018 | Gold et al.        | N/A | N/A         |
| 2019/0121889 | 12/2018 | Gold et al.        | N/A | N/A         |
| 2019/0158594 | 12/2018 | Shadmon et al.     | N/A | N/A         |
| 2019/0171438 | 12/2018 | Franchitti         | N/A | N/A         |
| 2019/0197246 | 12/2018 | Viswanathan et al. | N/A | N/A         |
| 2019/0244241 | 12/2018 | Hain               | N/A | N/A         |
| 2019/0245680 | 12/2018 | Boutaba et al.     | N/A | N/A         |
| 2019/0295079 | 12/2018 | Bae                | N/A | N/A         |
| 2019/0303483 | 12/2018 | Qian et al.        | N/A | N/A         |
| 2019/0303931 | 12/2018 | Valencia           | N/A | N/A         |
| 2019/0325517 | 12/2018 | Jayaram et al.     | N/A | N/A         |
| 2019/0354968 | 12/2018 | Sato et al.        | N/A | N/A         |
| 2019/0378352 | 12/2018 | Dey et al.         | N/A | N/A         |
| 2020/0099528 | 12/2019 | Chan               | N/A | N/A         |
| 2020/0117757 | 12/2019 | Yanamandra et al.  | N/A | N/A         |
| 2020/0151714 | 12/2019 | Chan et al.        | N/A | N/A         |
| 2020/0162266 | 12/2019 | Miller et al.      | N/A | N/A         |
| 2022/0156249 | 12/2021 | Luedtke            | N/A | N/A         |
| 2023/0062434 | 12/2022 | Wagner             | N/A | H04L 9/3247 |

## OTHER PUBLICATIONS

Bitincka, Ledion et al., “Optimizing Data Analysis with a Semi-structured Time Series Database,” self-published, first presented at “Workshop on Managing Systems via Log Analysis and Machine Learning Techniques (SLAML)”, Vancouver, British Columbia, Oct. 3, 2010. cited by applicant  
Carraso, David, “Exploring Splunk,” published by CITO Research, New York, NY, Apr. 2012. cited by applicant

SLAML 10 Reports, Workshop On Managing Systems via Log Analysis and Machine Learning Techniques, ;login: Feb. 2011 Conference Reports. cited by applicant

Splunk Enterprise 8.0.0 Overview, available online, retrieved May 20, 2020 from docs.splunk.com. cited by applicant

Splunk Cloud 8.0.2004 User Manual, available online, retrieved May 20, 2020 from docs.splunk.com. cited by applicant

Splunk Quick Reference Guide, updated 2019, available online at <https://www.splunk.com/pdfs/solution-guides/splunk-quick-reference-guide.pdf>, retrieved May 20, 2020. cited by applicant

---

*Primary Examiner:* Uddin; Md I

*Attorney, Agent or Firm:* Polsinelli P.C.

---

## Background/Summary

### BACKGROUND

(1) Information technology (IT) environments can include diverse types of data systems that store large amounts of diverse data types generated by numerous devices. For example, a big data ecosystem may include databases such as MySQL and Oracle databases, cloud computing services such as Amazon web services (AWS), and other data systems that store passively or actively generated data, including machine-generated data (“machine data”). The machine data can include

log data, performance data, diagnostic data, metrics, tracing data, or any other data that can be analyzed to diagnose equipment performance problems, monitor user interactions, and to derive other insights.

(2) The large amount and diversity of data systems containing large amounts of structured, semi-structured, and unstructured data relevant to any search query can be massive, and continues to grow rapidly. This technological evolution can give rise to various challenges in relation to managing, understanding and effectively utilizing the data. To reduce the potentially vast amount of data that may be generated, some data systems pre-process data based on anticipated data analysis needs. In particular, specified data items may be extracted from the generated data and stored in a data system to facilitate efficient retrieval and analysis of those data items at a later time. At least some of the remainder of the generated data is typically discarded during pre-processing.

(3) However, storing massive quantities of minimally processed or unprocessed data (collectively and individually referred to as “raw data”) for later retrieval and analysis is becoming increasingly more feasible as storage capacity becomes more inexpensive and plentiful. In general, storing raw data and performing analysis on that data later can provide greater flexibility because it enables an analyst to analyze all of the generated data instead of only a fraction of it. Although the availability of vastly greater amounts of diverse data on diverse data systems provides opportunities to derive new insights, it also gives rise to technical challenges to search and analyze the data in a performant way.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

(1) The examples described herein will be understood more fully from the detailed description given below and from the accompanying drawings, which, however, should not be taken to limit the application to the specific examples, but are for explanation and understanding.

(2) FIG. 1 is a block diagram of an example networked computer environment, in accordance with example embodiments;

(3) FIG. 2 is a block diagram of an example data intake and query system, in accordance with example embodiments;

(4) FIG. 3 is a block diagram of an example cloud-based data intake and query system, in accordance with example embodiments;

(5) FIG. 4 is a block diagram of an example data intake and query system that performs searches across external data systems, in accordance with example embodiments;

(6) FIG. 5A is a flowchart of an example method that illustrates how indexers process, index, and store data received from forwarders, in accordance with example embodiments;

(7) FIG. 5B is a block diagram of a data structure in which time-stamped event data can be stored in a data store, in accordance with example embodiments;

(8) FIG. 5C provides a visual representation of the manner in which a pipelined search language or query operates, in accordance with example embodiments;

(9) FIG. 6A is a flow diagram of an example method that illustrates how a search head and indexers perform a search query, in accordance with example embodiments;

(10) FIG. 6B provides a visual representation of an example manner in which a pipelined command language or query operates, in accordance with example embodiments;

(11) FIG. 7A is a diagram of an example scenario where a common customer identifier is found among log data received from three disparate data sources, in accordance with example embodiments;

(12) FIG. 7B illustrates an example of processing keyword searches and field searches, in accordance with disclosed embodiments;

- (13) FIG. 7C illustrates an example of creating and using an inverted index, in accordance with example embodiments;
- (14) FIG. 7D depicts a flowchart of example use of an inverted index in a pipelined search query, in accordance with example embodiments;
- (15) FIG. 8A is an interface diagram of an example user interface for a search screen, in accordance with example embodiments;
- (16) FIG. 8B is an interface diagram of an example user interface for a data summary dialog that enables a user to select various data sources, in accordance with example embodiments;
- (17) FIGS. 9, 10, 11A, 11B, 11C, 11D, 12, 13, 14, and 15 are interface diagrams of example report generation user interfaces, in accordance with example embodiments;
- (18) FIG. 16 is an example search query received from a client and executed by search peers, in accordance with example embodiments;
- (19) FIG. 17A is an interface diagram of an example user interface of a key indicators view, in accordance with example embodiments;
- (20) FIG. 17B is an interface diagram of an example user interface of an incident review dashboard, in accordance with example embodiments;
- (21) FIG. 17C is a tree diagram of an example a proactive monitoring tree, in accordance with example embodiments;
- (22) FIG. 17D is an interface diagram of an example a user interface displaying both log data and performance data, in accordance with example embodiments;
- (23) FIG. 18 is a block diagram of an embodiment of the data processing environment;
- (24) FIGS. 19A and 19B are block diagrams illustrating embodiments of a distributed ledger;
- (25) FIG. 20 is a block diagram illustrating an embodiment of a blockchain;
- (26) FIGS. 21A-21D are data flow diagrams illustrating an embodiment of a distributed ledger system processing a transaction;
- (27) FIG. 22 is a block diagram illustrating an embodiment of an environment that includes one or more nodes as a data source for the data intake and query system;
- (28) FIG. 23 is a data flow diagram illustrating transaction decoding by a distributed ledger connector of the data intake and query system 108 operating in accordance with aspects of the present disclosure;
- (29) FIG. 24 is a data flow diagram of computing a digital fingerprint of a smart contract, in accordance with aspects of the present disclosure;
- (30) FIG. 25 is a flow diagram of an embodiment of a method 2500 of decoding distributed ledger transactions, in accordance with aspects of the present disclosure; and
- (31) FIG. 26 is a flow diagram of an embodiment of a method 2600 of decoding transaction data, in accordance with aspects of the present disclosure.
- (32) FIG. 27 is a block diagram illustrating an embodiment of a data ingestion system.
- (33) FIG. 28 is a block diagram illustrating an embodiment of a data ingestion system.
- (34) FIG. 29 is a block diagram illustrating an embodiment of a data ingestion system.
- (35) FIG. 30 is an interface diagram of an example user interface, in accordance with example embodiments.
- (36) FIG. 31 is a flow diagram illustrative of an embodiment of a routine implemented by a processing system to route blockchain data to a data intake and query system.

#### DETAILED DESCRIPTION

(37) Described herein are systems and methods for decoding distributed ledger transaction records. The decoded transaction records may be utilized for data extraction and visualization, as well as for further processing, such as data analysis, alert triggering, and/or reporting, as described in more detail herein below.

(38) An example distributed ledger may include multiple nodes, such that each node may be associated with one or more distributed ledger accounts. The distributed ledger may implement a

transaction-based state machine, which transitions to a new state based on a set of inputs represented by transaction records (referred to as “transactions” for conciseness). A distributed ledger may be cryptographically-protected, e.g., by cryptographically encrypting the transaction records, such that reversing a transaction becomes computationally infeasible. In one embodiment, the cryptographically-protected distributed ledger may be implemented by a blockchain.

(39) A transaction may encode a message that is sent by a source account to a destination account. The message, which is signed by the private key of the source account, may specify a transfer of a certain amount of a digital asset from the source account to the destination account.

(40) Some distributed ledgers (e.g., Ethereum) support a special account type, which is referred to as “contract account.” A message to a contract account activates its executable code implementing a “smart contract,” which may evaluate specified conditions and perform various actions (e.g., transfer cryptocurrency tokens between accounts, write data to internal storage, mint new cryptocurrency tokens, perform calculation, create new smart contracts, etc.). The nodes of the distributed ledger may collectively implement a distributed virtual machine (e.g., the Ethereum Virtual Machine (EVM)) for executing the code implementing smart contracts. A smart contract can be created in a high level programming language (such as Solidity) and then compiled into the EVM bytecode.

(41) In various implementations, a transaction may further specify various other parameters, e.g., the amount of a digital asset to be transferred to a node that has successfully processed the transaction. A transaction may be cryptographically signed by the originating node.

(42) To cause a state transition of the blockchain, a transaction should be validated by at least one node, which would then include it, together with other transactions, into a block that is appended to the blockchain. The block also includes a “proof of work” value that has been computed by the node that created the block in order to enforce a sequential order of blocks. The proof of work value is produced by solving a computationally intensive task (e.g., computing padding bits to be appended to the block in order to produce a predetermined value of the block hash). Some distributed ledgers, such as Ethereum, may rely on other consensus mechanisms, such as Clique, IBFT, PBFT, or use proof of stake instead of proof of work to order blocks. Thus, a distributed ledger may implement a cryptographically protected distributed immutable database and a distributed virtual machine for executing smart contracts.

(43) A data intake and query system operating in accordance with aspects of the present disclosure can implement a Getting-Data-In (GDI) component (such as data adapter, monitor, forwarder, connector, or the like) in order to ingest the distributed ledger transaction data, e.g., by reading log files maintained by one or more nodes of the distributed ledger, listening to the blocks, transactions, and events that are broadcasted to all participating nodes of a distributed ledger, and/or performing other actions. The ingested raw data may be aggregated, decoded, visualized, and/or further processed by the data intake and query system.

(44) As noted herein above, a transaction can include invocation of a smart contract. Since the bytecode that encodes smart contract invocation does not preserve parameter names, meaningful decoding of such a transaction would require extrinsic knowledge of the meanings of the parameters. Decoding the transaction may be further hindered by potentially overlapping signatures of functions exposed by different smart contracts, i.e., two or more different functions having the same signature that encodes the function name and parameter types. In an illustrative example, each of two smart contracts may expose a function with the signature transfer (address, uint256). In one smart contract, the second argument (having the type of uint256) may refer to an amount of digital currency being transferred, while in another smart contract the second argument may refer to a token identifier. However, both functions would have the same signature, thus presenting a challenge for the transaction decoder.

(45) The systems and methods of the present disclosure overcome the above-noted challenges by implementing digital fingerprinting of smart contracts, which facilitates associating a smart

contract of interest with a known application binary interface (ABI) definition. “Digital fingerprint” herein refers to a numeric value (represented by a bit string of a predetermined size) which can be unambiguously derived from the smart contract bytecode encoding or ABI definition, such that the probability of two different smart contracts (e.g., smart contracts that have different ABI definitions) having the same digital fingerprint value is very low (e.g., below a predetermined probability threshold).

(46) In an illustrative example, the digital fingerprint of a smart contract can be represented by a hash of all function and event signatures exposed by the contract ABI. Thus, a transaction decoder implemented in accordance with aspects of the present disclosure may compute a digital fingerprint for the EVM bytecode implementing a smart contract invoked by a distributed ledger transaction and compare the computed digital fingerprint with known digital fingerprints of ABI definitions of smart contracts, the database of which can be maintained by the data intake and query system. Upon identifying a matching smart contract among the stored ABI definitions, signatures of the functions and events of the identified matching smart contracts, along with other pertinent information that can be stored in the smart contract ABI database maintained by the query intake and analysis system, can be utilized for decoding the distributed ledger transaction. The decoded transaction data, including the function name, the parameters names, types, and values, may be fed to a data intake and query system, such as the SPLUNK® ENTERPRISE system developed by Splunk Inc. of San Francisco, California, as described in more detail herein below.

(47) The SPLUNK® ENTERPRISE system is the leading platform for providing real-time operational intelligence that enables organizations to collect, index, and search machine data from various websites, applications, servers, networks, and mobile devices that power their businesses. The data intake and query system is particularly useful for analyzing data which is commonly found in system log files, network data, and other data input sources. Although many of the techniques described herein are explained with reference to a data intake and query system similar to the SPLUNK® ENTERPRISE system, these techniques are also applicable to other types of data systems.

(48) In the data intake and query system, machine data are collected and stored as “events.” An event comprises a portion of machine data and is associated with a specific point in time. The portion of machine data may reflect activity in an IT environment and may be produced by a component of that IT environment, where the events may be searched to provide insight into the IT environment, thereby improving the performance of components in the IT environment. Events may be derived from “time series data,” where the time series data comprises a sequence of data points (e.g., performance measurements from a computer system, etc.) that are associated with successive points in time. In general, each event has a portion of machine data that is associated with a timestamp that is derived from the portion of machine data in the event. A timestamp of an event may be determined through interpolation between temporally proximate events having known timestamps or may be determined based on other configurable rules for associating timestamps with events.

(49) In some instances, machine data can have a predefined format, where data items with specific data formats are stored at predefined locations in the data. For example, the machine data may include data associated with fields in a database table. In other instances, machine data may not have a predefined format (e.g., may not be at fixed, predefined locations), but may have repeatable (e.g., non-random) patterns. This means that some machine data can comprise various data items of different data types that may be stored at different locations within the data. For example, when the data source is an operating system log, an event can include one or more lines from the operating system log containing machine data that includes different types of performance and diagnostic information associated with a specific point in time (e.g., a timestamp).

(50) Examples of components which may generate machine data from which events can be derived include, but are not limited to, web servers, application servers, databases, firewalls, routers,

operating systems, and software applications that execute on computer systems, mobile devices, sensors, Internet of Things (IoT) devices, distributed ledger nodes, etc. The machine data generated by such data sources can include, for example and without limitation, server log files, activity log files, configuration files, messages, network packet data, performance measurements, sensor measurements, distributed ledger transactions, etc.

(51) The data intake and query system uses a flexible schema to specify how to extract information from events. A flexible schema may be developed and redefined as needed. Note that a flexible schema may be applied to events “on the fly,” when it is needed (e.g., at search time, index time, ingestion time, etc.). When the schema is not applied to events until search time, the schema may be referred to as a “late-binding schema.”

(52) During operation, the data intake and query system receives machine data from any type and number of sources (e.g., one or more system logs, streams of network packet data, sensor data, application program data, error logs, stack traces, system performance data, etc.). The system parses the machine data to produce events each having a portion of machine data associated with a timestamp. The system stores the events in a data store. The system enables users to run queries against the stored events to, for example, retrieve events that meet criteria specified in a query, such as criteria indicating certain keywords or having specific values in defined fields. As used herein, the term “field” refers to a location in the machine data of an event containing one or more values for a specific data item. A field may be referenced by a field name associated with the field. As will be described in more detail herein, a field is defined by an extraction rule (e.g., a regular expression) that derives one or more values or a sub-portion of text from the portion of machine data in each event to produce a value for the field for that event. The set of values produced are semantically-related (such as IP address), even though the machine data in each event may be in different formats (e.g., semantically-related values may be in different positions in the events derived from different sources).

(53) As described above, the system stores the events in a data store. The events stored in the data store are field-searchable, where field-searchable herein refers to the ability to search the machine data (e.g., the raw machine data) of an event based on a field specified in search criteria. For example, a search having criteria that specifies a field name “UserID” may cause the system to field-search the machine data of events to identify events that have the field name “UserID.” In another example, a search having criteria that specifies a field name “UserID” with a corresponding field value “12345” may cause the system to field-search the machine data of events to identify events having that field-value pair (e.g., field name “UserID” with a corresponding field value of “12345”). Events are field-searchable using one or more configuration files associated with the events. Each configuration file includes one or more field names, where each field name is associated with a corresponding extraction rule and a set of events to which that extraction rule applies. The set of events to which an extraction rule applies may be identified by metadata associated with the set of events. For example, an extraction rule may apply to a set of events that are each associated with a particular host, source, or source type. When events are to be searched based on a particular field name specified in a search, the system uses one or more configuration files to determine whether there is an extraction rule for that particular field name that applies to each event that falls within the criteria of the search. If so, the event is considered as part of the search results (and additional processing may be performed on that event based on criteria specified in the search). If not, the next event is similarly analyzed, and so on.

(54) As noted above, the data intake and query system utilizes a late-binding schema while performing queries on events. One aspect of a late-binding schema is applying extraction rules to events to extract values for specific fields during search time. More specifically, the extraction rule for a field can include one or more instructions that specify how to extract a value for the field from an event. An extraction rule can generally include any type of instruction for extracting values from events. In some cases, an extraction rule comprises a regular expression, where a sequence of



characters forms a search pattern. An extraction rule comprising a regular expression is referred to herein as a regex rule. The system applies a regex rule to an event to extract values for a field associated with the regex rule, where the values are extracted by searching the event for the sequence of characters defined in the regex rule.

(55) In the data intake and query system, a field extractor may be configured to automatically generate extraction rules for certain fields in the events when the events are being created, indexed, or stored, or possibly at a later time. Alternatively, a user may manually define extraction rules for fields using a variety of techniques. In contrast to a conventional schema for a database system, a late-binding schema is not defined at data ingestion time. Instead, the late-binding schema can be developed on an ongoing basis until the time a query is actually executed. This means that extraction rules for the fields specified in a query may be provided in the query itself, or may be located during execution of the query. Hence, as a user learns more about the data in the events, the user can continue to refine the late-binding schema by adding new fields, deleting fields, or modifying the field extraction rules for use the next time the schema is used by the system. Because the data intake and query system maintains the underlying machine data and uses a late-binding schema for searching the machine data, it enables a user to continue investigating and learn valuable insights about the machine data.

(56) In some embodiments, a common field name may be used to reference two or more fields containing equivalent and/or similar data items, even though the fields may be associated with different types of events that possibly have different data formats and different extraction rules. By enabling a common field name to be used to identify equivalent and/or similar fields from different types of events generated by disparate data sources, the system facilitates use of a “common information model” (CIM) across the disparate data sources (further discussed with respect to FIG. 7A).

(57) FIG. 1 is a block diagram of an example networked computer environment **100**, in accordance with example embodiments. Those skilled in the art would understand that FIG. 1 represents one example of a networked computer system and other embodiments may use different arrangements.

(58) The networked computer system **100** comprises one or more computing devices. These one or more computing devices comprise any combination of hardware and software configured to implement the various logical components described herein. For example, the one or more computing devices may include one or more memories that store instructions for implementing the various components described herein, one or more hardware processors configured to execute the instructions stored in the one or more memories, and various data repositories in the one or more memories for storing data structures utilized and manipulated by the various components.

(59) In some embodiments, one or more client devices **102** are coupled to one or more host devices **106** and a data intake and query system **108** via one or more networks **104**. Networks **104** broadly represent one or more LANs, WANs, cellular networks (e.g., LTE, HSPA, 3G, and other cellular technologies), and/or networks using any of wired, wireless, terrestrial microwave, or satellite links, and may include the public Internet.

(60) In the illustrated embodiment, a system **100** includes one or more host devices **106**. Host devices **106** may broadly include any number of computers, virtual machine instances, and/or data centers that are configured to host or execute one or more instances of host applications **114**. In general, a host device **106** may be involved, directly or indirectly, in processing requests received from client devices **102**. Each host device **106** may comprise, for example, one or more of a network device, a web server, an application server, a database server, etc. A collection of host devices **106** may be configured to implement a network-based service. For example, a provider of a network-based service may configure one or more host devices **106** and host applications **114** (e.g., one or more web servers, application servers, database servers, etc.) to collectively implement the network-based application.

(61) In general, client devices **102** communicate with one or more host applications **114** to

exchange information. The communication between a client device **102** and a host application **114** may, for example, be based on the Hypertext Transfer Protocol (HTTP) or any other network protocol. Content delivered from the host application **114** to a client device **102** may include, for example, HTML documents, media content, etc. The communication between a client device **102** and host application **114** may include sending various requests and receiving data packets. For example, in general, a client device **102** or application running on a client device may initiate communication with a host application **114** by making a request for a specific resource (e.g., based on an HTTP request), and the application server may respond with the requested content stored in one or more response packets.

(62) In the illustrated embodiment, one or more of host applications **114** may generate various types of performance data during operation, including event logs, network data, sensor data, and other types of machine data. For example, a host application **114** comprising a web server may generate one or more web server logs in which details of interactions between the web server and any number of client devices **102** is recorded. As another example, a host device **106** comprising a router may generate one or more router logs that record information related to network traffic managed by the router. As yet another example, a host application **114** comprising a database server may generate one or more logs that record information related to requests sent from other host applications **114** (e.g., web servers or application servers) for data managed by the database server.

(63) Client devices **102** of FIG. **1** represent any computing device capable of interacting with one or more host devices **106** via a network **104**. Examples of client devices **102** may include, without limitation, smart phones, tablet computers, handheld computers, wearable devices, laptop computers, desktop computers, servers, portable media players, gaming devices, and so forth. In general, a client device **102** can provide access to different content, for instance, content provided by one or more host devices **106**, etc. Each client device **102** may comprise one or more client applications **110**, described in more detail in a separate section hereinafter.

(64) In some embodiments, each client device **102** may host or execute one or more client applications **110** that are capable of interacting with one or more host devices **106** via one or more networks **104**. For instance, a client application **110** may be or comprise a web browser that a user may use to navigate to one or more websites or other resources provided by one or more host devices **106**. As another example, a client application **110** may comprise a mobile application or “app.” For example, an operator of a network-based service hosted by one or more host devices **106** may make available one or more mobile apps that enable users of client devices **102** to access various resources of the network-based service. As yet another example, client applications **110** may include background processes that perform various operations without direct interaction from a user. A client application **110** may include a “plug-in” or “extension” to another application, such as a web browser plug-in or extension.

(65) In some embodiments, a client application **110** may include a monitoring component **112**. At a high level, the monitoring component **112** comprises a software component or other logic that facilitates generating performance data related to a client device's operating state, including monitoring network traffic sent and received from the client device and collecting other device and/or application-specific information. Monitoring component **112** may be an integrated component of a client application **110**, a plug-in, an extension, or any other type of add-on component. Monitoring component **112** may also be a stand-alone process.

(66) In some embodiments, a monitoring component **112** may be created when a client application **110** is developed, for example, by an application developer using a software development kit (SDK). The SDK may include custom monitoring code that can be incorporated into the code implementing a client application **110**. When the code is converted to an executable application, the custom code implementing the monitoring functionality can become part of the application itself.

(67) In some embodiments, an SDK or other code for implementing the monitoring functionality may be offered by a provider of a data intake and query system, such as a system **108**. In such

cases, the provider of the system **108** can implement the custom code so that performance data generated by the monitoring functionality is sent to the system **108** to facilitate analysis of the performance data by a developer of the client application or other users.

(68) In some embodiments, the custom monitoring code may be incorporated into the code of a client application **110** in a number of different ways, such as the insertion of one or more lines in the client application code that call or otherwise invoke the monitoring component **112**. As such, a developer of a client application **110** can add one or more lines of code into the client application **110** to trigger the monitoring component **112** at desired points during execution of the application. Code that triggers the monitoring component may be referred to as a monitor trigger. For instance, a monitor trigger may be included at or near the beginning of the executable code of the client application **110** such that the monitoring component **112** is initiated or triggered as the application is launched, or included at other points in the code that correspond to various actions of the client application, such as sending a network request or displaying a particular interface.

(69) In some embodiments, the monitoring component **112** may monitor one or more aspects of network traffic sent and/or received by a client application **110**. For example, the monitoring component **112** may be configured to monitor data packets transmitted to and/or from one or more host applications **114**. Incoming and/or outgoing data packets can be read or examined to identify network data contained within the packets, for example, and other aspects of data packets can be analyzed to determine a number of network performance statistics. Monitoring network traffic may enable information to be gathered particular to the network performance associated with a client application **110** or set of applications.

(70) In some embodiments, network performance data refers to any type of data that indicates information about the network and/or network performance. Network performance data may include, for instance, a URL requested, a connection type (e.g., HTTP, HTTPS, etc.), a connection start time, a connection end time, an HTTP status code, request length, response length, request headers, response headers, connection status (e.g., completion, response time(s), failure, etc.), and the like. Upon obtaining network performance data indicating performance of the network, the network performance data can be transmitted to a data intake and query system **108** for analysis.

(71) Upon developing a client application **110** that incorporates a monitoring component **112**, the client application **110** can be distributed to client devices **102**. Applications generally can be distributed to client devices **102** in any manner, or they can be pre-loaded. In some cases, the application may be distributed to a client device **102** via an application marketplace or other application distribution system. For instance, an application marketplace or other application distribution system might distribute the application to a client device based on a request from the client device to download the application.

(72) Examples of functionality that enables monitoring performance of a client device are described in U.S. patent application Ser. No. 14/524,748, entitled “UTILIZING PACKET HEADERS TO MONITOR NETWORK TRAFFIC IN ASSOCIATION WITH A CLIENT DEVICE,” filed on 27 Oct. 2014, and which is hereby incorporated by reference in its entirety for all purposes.

(73) In some embodiments, the monitoring component **112** may also monitor and collect performance data related to one or more aspects of the operational state of a client application **110** and/or client device **102**. For example, a monitoring component **112** may be configured to collect device performance information by monitoring one or more client device operations, or by making calls to an operating system and/or one or more other applications executing on a client device **102** for performance information. Device performance information may include, for instance, a current wireless signal strength of the device, a current connection type and network carrier, current memory performance information, a geographic location of the device, a device orientation, and any other information related to the operational state of the client device.

(74) In some embodiments, the monitoring component **112** may also monitor and collect other

device profile information including, for example, a type of client device, a manufacturer and model of the device, versions of various software applications installed on the device, and so forth. (75) In general, a monitoring component **112** may be configured to generate performance data in response to a monitor trigger in the code of a client application **110** or other triggering application event, as described above, and to store the performance data in one or more data records. Each data record, for example, may include a collection of field-value pairs, each field-value pair storing a particular item of performance data in association with a field for the item. For example, a data record generated by a monitoring component **112** may include a “network Latency” field (not shown in the Figure) in which a value is stored. This field indicates a network latency measurement associated with one or more network requests. The data record may include a “state” field to store a value indicating a state of a network connection, and so forth for any number of aspects of collected performance data.

(76) FIG. 2 is a block diagram of an example data intake and query system **108**, in accordance with example embodiments. System **108** includes one or more forwarders **204** that receive data from a variety of input data sources **202**, and one or more indexers **206** that process and store the data in one or more data stores **208**. These forwarders **204** and indexers **206** can comprise separate computer systems, or may alternatively comprise separate processes executing on one or more computer systems.

(77) Each data source **202** broadly represents a distinct source of data that can be consumed by system **108**. Examples of a data sources **202** include, without limitation, data files, directories of files, data sent over a network, event logs, registries, etc.

(78) During operation, the forwarders **204** identify which indexers **206** receive data collected from a data source **202** and forward the data to the appropriate indexers. Forwarders **204** can also perform operations on the data before forwarding, including removing extraneous data, detecting timestamps in the data, parsing data, indexing data, routing data based on criteria relating to the data being routed, and/or performing other data transformations.

(79) In some embodiments, a forwarder **204** may comprise a service accessible to client devices **102** and host devices **106** via a network **104**. For example, one type of forwarder **204** may be capable of consuming vast amounts of real-time data from a potentially large number of client devices **102** and/or host devices **106**. The forwarder **204** may, for example, comprise a computing device which implements multiple data pipelines or “queues” to handle forwarding of network data to indexers **206**. A forwarder **204** may also perform many of the functions that are performed by an indexer. For example, a forwarder **204** may perform keyword extractions on raw data or parse raw data to create events. A forwarder **204** may generate time stamps for events. Additionally, or alternatively, a forwarder **204** may perform routing of events to indexers **206**. Data store **208** may contain events derived from machine data from a variety of sources all pertaining to the same component in an IT environment, and this data may be produced by the machine in question or by other components in the IT environment.

(80) The example data intake and query system **108** described in reference to FIG. 2 comprises several system components, including one or more forwarders, indexers, and search heads. In some environments, a user of a data intake and query system **108** may install and configure, on computing devices owned and operated by the user, one or more software applications that implement some or all of these system components. For example, a user may install a software application on server computers owned by the user and configure each server to operate as one or more of a forwarder, an indexer, a search head, etc. This arrangement generally may be referred to as an “on-premises” solution. That is, the system **108** is installed and operates on computing devices directly controlled by the user of the system. Some users may prefer an on-premises solution because it may provide a greater level of control over the configuration of certain aspects of the system (e.g., security, privacy, standards, controls, etc.). However, other users may instead prefer an arrangement in which the user is not directly responsible for providing and managing the

computing devices upon which various components of system **108** operate.

(81) In one embodiment, to provide an alternative to an entirely on-premises environment for system **108**, one or more of the components of a data intake and query system instead may be provided as a cloud-based service. In this context, a cloud-based service refers to a service hosted by one or more computing resources that are accessible to end users over a network, for example, by using a web browser or other application on a client device to interface with the remote computing resources. For example, a service provider may provide a cloud-based data intake and query system by managing computing resources configured to implement various aspects of the system (e.g., forwarders, indexers, search heads, etc.) and by providing access to the system to end users via a network. Typically, a user may pay a subscription or other fee to use such a service. Each subscribing user of the cloud-based service may be provided with an account that enables the user to configure a customized cloud-based system based on the user's preferences.

(82) FIG. 3 illustrates a block diagram of an example cloud-based data intake and query system. Similar to the system of FIG. 2, the networked computer system **300** includes input data sources **202** and forwarders **204**. These input data sources and forwarders may be in a subscriber's private computing environment. Alternatively, they might be directly managed by the service provider as part of the cloud service. In the example system **300**, one or more forwarders **204** and client devices **302** are coupled to a cloud-based data intake and query system **306** via one or more networks **304**. Network **304** broadly represents one or more LANs, WANs, cellular networks, intranetworks, internetworks, etc., using any of wired, wireless, terrestrial microwave, satellite links, etc., and may include the public Internet, and is used by client devices **302** and forwarders **204** to access the system **306**. Similar to the system of **38**, each of the forwarders **204** may be configured to receive data from an input source and to forward the data to other components of the system **306** for further processing.

(83) In some embodiments, a cloud-based data intake and query system **306** may comprise a plurality of system instances **308**. In general, each system instance **308** may include one or more computing resources managed by a provider of the cloud-based system **306** made available to a particular subscriber. The computing resources comprising a system instance **308** may, for example, include one or more servers or other devices configured to implement one or more forwarders, indexers, search heads, and other components of a data intake and query system, similar to system **108**. As indicated above, a subscriber may use a web browser or other application of a client device **302** to access a web portal or other interface that enables the subscriber to configure an instance **308**.

(84) Providing a data intake and query system as described in reference to system **108** as a cloud-based service presents a number of challenges. Each of the components of a system **108** (e.g., forwarders, indexers, and search heads) may at times refer to various configuration files stored locally at each component. These configuration files typically may involve some level of user configuration to accommodate particular types of data a user desires to analyze and to account for other user preferences. However, in a cloud-based service context, users typically may not have direct access to the underlying computing resources implementing the various system components (e.g., the computing resources comprising each system instance **308**) and may desire to make such configurations indirectly, for example, using one or more web-based interfaces. Thus, the techniques and systems described herein for providing user interfaces that enable a user to configure source type definitions are applicable to both on-premises and cloud-based service contexts, or some combination thereof (e.g., a hybrid system where both an on-premises environment, such as SPLUNK® ENTERPRISE, and a cloud-based environment, such as SPLUNK CLOUD□, are centrally visible).

(85) FIG. 4 shows a block diagram of an example of a data intake and query system **108** that provides transparent search facilities for data systems that are external to the data intake and query system. Such facilities are available in the Splunk® Analytics for Hadoop® system provided by

Splunk Inc. of San Francisco, California. Splunk® Analytics for Hadoop® represents an analytics platform that enables business and IT teams to rapidly explore, analyze, and visualize data in Hadoop® and NoSQL data stores.

(86) The search head **210** of the data intake and query system receives search requests from one or more client devices **404** over network connections **420**. As discussed above, the data intake and query system **108** may reside in an enterprise location, in the cloud, etc. FIG. 4 illustrates that multiple client devices **404a**, **404b** . . . **404n** may communicate with the data intake and query system **108**. The client devices **404** may communicate with the data intake and query system using a variety of connections. For example, one client device in FIG. 4 is illustrated as communicating over an Internet (Web) protocol, another client device is illustrated as communicating via a command line interface, and another client device is illustrated as communicating via a software developer kit (SDK).

(87) The search head **210** analyzes the received search request to identify request parameters. If a search request received from one of the client devices **404** references an index maintained by the data intake and query system, then the search head **210** connects to one or more indexers **206** of the data intake and query system for the index referenced in the request parameters. That is, if the request parameters of the search request reference an index, then the search head accesses the data in the index via the indexer. The data intake and query system **108** may include one or more indexers **206**, depending on system access resources and requirements. As described further below, the indexers **206** retrieve data from their respective local data stores **208** as specified in the search request. The indexers and their respective data stores can comprise one or more storage devices and typically reside on the same system, though they may be connected via a local network connection.

(88) If the request parameters of the received search request reference an external data collection, which is not accessible to the indexers **206** or under the management of the data intake and query system, then the search head **210** can access the external data collection through an External Result Provider (ERP) process **410**. An external data collection may be referred to as a “virtual index” (plural, “virtual indices”). An ERP process provides an interface through which the search head **210** may access virtual indices.

(89) Thus, a search reference to an index of the system relates to a locally stored and managed data collection. In contrast, a search reference to a virtual index relates to an externally stored and managed data collection, which the search head may access through one or more ERP processes **410**, **412**. FIG. 4 shows two ERP processes **410**, **412** that connect to respective remote (external) virtual indices, which are indicated as a Hadoop or another system **414** (e.g., Amazon S3, Amazon EMR, other Hadoop® Compatible File Systems (HCFS), etc.) and a relational database management system (RDBMS) **416**. Other virtual indices may include other file organizations and protocols, such as Structured Query Language (SQL) and the like. The ellipses between the ERP processes **410**, **412** indicate optional additional ERP processes of the data intake and query system **108**. An ERP process may be a computer process that is initiated or spawned by the search head **210** and is executed by the search data intake and query system **108**. Alternatively, or additionally, an ERP process may be process spawned by the search head **210** on the same or different host system as the search head **210** resides.

(90) The search head **210** may spawn a single ERP process in response to multiple virtual indices referenced in a search request, or the search head may spawn different ERP processes for different virtual indices. Generally, virtual indices that share common data configurations or protocols may share ERP processes. For example, all search query references to a Hadoop file system may be processed by the same ERP process, if the ERP process is suitably configured. Likewise, all search query references to a SQL database may be processed by the same ERP process. In addition, the search head may provide a common ERP process for common external data source types (e.g., a common vendor may utilize a common ERP process, even if the vendor includes different data storage system types, such as Hadoop and SQL). Common indexing schemes also may be handled

by common ERP processes, such as flat text files or Weblog files.

(91) The search head **210** determines the number of ERP processes to be initiated via the use of configuration parameters that are included in a search request message. Generally, there is a one-to-many relationship between an external results provider “family” and ERP processes. There is also a one-to-many relationship between an ERP process and corresponding virtual indices that are referred to in a search request. For example, using RDBMS, assume two independent instances of such a system by one vendor, such as one RDBMS for production and another RDBMS used for development. In such a situation, it is likely preferable (but optional) to use two ERP processes to maintain the independent operation as between production and development data. Both of the ERPs, however, will belong to the same family, because the two RDBMS system types are from the same vendor.

(92) The ERP processes **410, 412** receive a search request from the search head **210**. The search head may optimize the received search request for execution at the respective external virtual index. Alternatively, the ERP process may receive a search request as a result of analysis performed by the search head or by a different system process. The ERP processes **410, 412** can communicate with the search head **210** via conventional input/output routines (e.g., standard in/standard out, etc.). In this way, the ERP process receives the search request from a client device such that the search request may be efficiently executed at the corresponding external virtual index.

(93) The ERP processes **410, 412** may be implemented as a process of the data intake and query system. Each ERP process may be provided by the data intake and query system, or may be provided by process or application providers who are independent of the data intake and query system. Each respective ERP process may include an interface application installed at a computer of the external result provider that ensures proper communication between the search support system and the external result provider. The ERP processes **410, 412** generate appropriate search requests in the protocol and syntax of the respective virtual indices **414, 416**, each of which corresponds to the search request received by the search head **210**. Upon receiving search results from their corresponding virtual indices, the respective ERP process passes the result to the search head **210**, which may return or display the results or a processed set of results based on the returned results to the respective client device.

(94) Client devices **404** may communicate with the data intake and query system **108** through a network interface **420**, e.g., one or more LANs, WANs, cellular networks, intranetworks, and/or internetworks using any of wired, wireless, terrestrial microwave, satellite links, etc., and may include the public Internet.

(95) The analytics platform utilizing the External Result Provider process described in more detail in U.S. Pat. No. 8,738,629, entitled “EXTERNAL RESULT PROVIDED PROCESS FOR RETRIEVING DATA STORED USING A DIFFERENT CONFIGURATION OR PROTOCOL”, issued on 27 May 2014, U.S. Pat. No. 8,738,587, entitled “PROCESSING A SYSTEM SEARCH REQUEST BY RETRIEVING RESULTS FROM BOTH A NATIVE INDEX AND A VIRTUAL INDEX”, issued on 25 Jul. 2013, U.S. patent application Ser. No. 14/266,832, entitled “PROCESSING A SYSTEM SEARCH REQUEST ACROSS DISPARATE DATA COLLECTION SYSTEMS”, filed on 1 May 2014, and U.S. Pat. No. 9,514,189, entitled “PROCESSING A SYSTEM SEARCH REQUEST INCLUDING EXTERNAL DATA SOURCES”, issued on 6 Dec. 2016, each of which is hereby incorporated by reference in its entirety for all purposes.

(96) The ERP processes described above may include two operation modes: a streaming mode and a reporting mode. The ERP processes can operate in streaming mode only, in reporting mode only, or in both modes simultaneously. Operating in both modes simultaneously is referred to as mixed mode operation. In a mixed mode operation, the ERP at some point can stop providing the search head with streaming results and only provide reporting results thereafter, or the search head at some point may start ignoring streaming results it has been using and only use reporting results thereafter.

(97) The streaming mode returns search results in real time, with minimal processing, in response to the search request. The reporting mode provides results of a search request with processing of the search results prior to providing them to the requesting search head, which in turn provides results to the requesting client device. ERP operation with such multiple modes provides greater performance flexibility with regard to report time, search latency, and resource utilization.

(98) In a mixed mode operation, both streaming mode and reporting mode are operating simultaneously. The streaming mode results (e.g., the machine data obtained from the external data source) are provided to the search head, which can then process the results data (e.g., break the machine data into events, timestamp it, filter it, etc.) and integrate the results data with the results data from other external data sources, and/or from data stores of the search head. The search head performs such processing and can immediately start returning interim (streaming mode) results to the user at the requesting client device; simultaneously, the search head is waiting for the ERP process to process the data it is retrieving from the external data source as a result of the concurrently executing reporting mode.

(99) In some instances, the ERP process initially operates in a mixed mode, such that the streaming mode operates to enable the ERP quickly to return interim results (e.g., some of the machined data or unprocessed data necessary to respond to a search request) to the search head, enabling the search head to process the interim results and begin providing to the client or search requester interim results that are responsive to the query. Meanwhile, in this mixed mode, the ERP also operates concurrently in reporting mode, processing portions of machine data in a manner responsive to the search query. Upon determining that it has results from the reporting mode available to return to the search head, the ERP may halt processing in the mixed mode at that time (or some later time) by stopping the return of data in streaming mode to the search head and switching to reporting mode only. The ERP at this point starts sending interim results in reporting mode to the search head, which in turn may then present this processed data responsive to the search request to the client or search requester. Typically, the search head switches from using results from the ERP's streaming mode of operation to results from the ERP's reporting mode of operation when the higher bandwidth results from the reporting mode outstrip the amount of data processed by the search head in the streaming mode of ERP operation.

(100) A reporting mode may have a higher bandwidth because the ERP does not have to spend time transferring data to the search head for processing all the machine data. In addition, the ERP may optionally direct another processor to do the processing.

(101) The streaming mode of operation does not need to be stopped to gain the higher bandwidth benefits of a reporting mode; the search head could simply stop using the streaming mode results—and start using the reporting mode results—when the bandwidth of the reporting mode has caught up with or exceeded the amount of bandwidth provided by the streaming mode. Thus, a variety of triggers and ways to accomplish a search head's switch from using streaming mode results to using reporting mode results may be appreciated by one skilled in the art.

(102) The reporting mode can involve the ERP process (or an external system) performing event breaking, time stamping, filtering of events to match the search query request, and calculating statistics on the results. The user can request particular types of data, such as if the search query itself involves types of events, or the search request may ask for statistics on data, such as on events that meet the search request. In either case, the search head understands the query language used in the received query request, which may be a proprietary language. One exemplary query language is Splunk Processing Language (SPL) developed by the assignee of the application, Splunk Inc. The search head typically understands how to use that language to obtain data from the indexers, which store data in a format used by the SPLUNK® Enterprise system.

(103) The ERP processes support the search head, as the search head is not ordinarily configured to understand the format in which data is stored in external data sources such as Hadoop or SQL data systems. Rather, the ERP process performs that translation from the query submitted in the search



support system's native format (e.g., SPL if SPLUNK® ENTERPRISE is used as the search support system) to a search query request format that will be accepted by the corresponding external data system. The external data system typically stores data in a different format from that of the search support system's native index format, and it utilizes a different query language (e.g., SQL or MapReduce, rather than SPL or the like).

(104) As noted, the ERP process can operate in the streaming mode alone. After the ERP process has performed the translation of the query request and received raw results from the streaming mode, the search head can integrate the returned data with any data obtained from local data sources (e.g., native to the search support system), other external data sources, and other ERP processes (if such operations were required to satisfy the terms of the search query). An advantage of mixed mode operation is that, in addition to streaming mode, the ERP process is also executing concurrently in reporting mode. Thus, the ERP process (rather than the search head) is processing query results (e.g., performing event breaking, timestamping, filtering, possibly calculating statistics if required to be responsive to the search query request, etc.). It should be apparent to those skilled in the art that additional time is needed for the ERP process to perform the processing in such a configuration. Therefore, the streaming mode will allow the search head to start returning interim results to the user at the client device before the ERP process can complete sufficient processing to start returning any search results. The switchover between streaming and reporting mode happens when the ERP process determines that the switchover is appropriate, such as when the ERP process determines it can begin returning meaningful results from its reporting mode.

(105) The operation described above illustrates the source of operational latency: streaming mode has low latency (immediate results) and usually has relatively low bandwidth (fewer results can be returned per unit of time). In contrast, the concurrently running reporting mode has relatively high latency (it has to perform a lot more processing before returning any results) and usually has relatively high bandwidth (more results can be processed per unit of time). For example, when the ERP process does begin returning report results, it returns more processed results than in the streaming mode, because, e.g., statistics only need to be calculated to be responsive to the search request. That is, the ERP process doesn't have to take time to first return machine data to the search head. As noted, the ERP process could be configured to operate in streaming mode alone and return just the machine data for the search head to process in a way that is responsive to the search request. Alternatively, the ERP process can be configured to operate in the reporting mode only. Also, the ERP process can be configured to operate in streaming mode and reporting mode concurrently, as described, with the ERP process stopping the transmission of streaming results to the search head when the concurrently running reporting mode has caught up and started providing results. The reporting mode does not require the processing of all machine data that is responsive to the search query request before the ERP process starts returning results; rather, the reporting mode usually performs processing of chunks of events and returns the processing results to the search head for each chunk.

(106) For example, an ERP process can be configured to merely return the contents of a search result file verbatim, with little or no processing of results. That way, the search head performs all processing (such as parsing byte streams into events, filtering, etc.). The ERP process can be configured to perform additional intelligence, such as analyzing the search request and handling all the computation that a native search indexer process would otherwise perform. In this way, the configured ERP process provides greater flexibility in features while operating according to desired preferences, such as response latency and resource requirements.

(107) FIG. 5A is a flow chart of an example method that illustrates how indexers process, index, and store data received from forwarders, in accordance with example embodiments. The data flow illustrated in FIG. 5A is provided for illustrative purposes only; those skilled in the art would understand that one or more of the steps of the processes illustrated in FIG. 5A may be removed or that the ordering of the steps may be changed. Furthermore, for the purposes of illustrating a clear

example, one or more particular system components are described in the context of performing various operations during each of the data flow stages. For example, a forwarder is described as receiving and processing machine data during an input phase; an indexer is described as parsing and indexing machine data during parsing and indexing phases; and a search head is described as performing a search query during a search phase. However, other system arrangements and distributions of the processing steps across system components may be used.

(108) At block **502**, a forwarder receives data from an input source, such as a data source **202** shown in FIG. 2. A forwarder initially may receive the data as a raw data stream generated by the input source. For example, a forwarder may receive a data stream from a log file generated by an application server, from a stream of network data from a network device, or from any other source of data. In some embodiments, a forwarder receives the raw data and may segment the data stream into “blocks,” possibly of a uniform data size, to facilitate subsequent processing steps.

(109) At block **504**, a forwarder or other system component annotates each block generated from the raw data with one or more metadata fields. These metadata fields may, for example, provide information related to the data block as a whole and may apply to each event that is subsequently derived from the data in the data block. For example, the metadata fields may include separate fields specifying each of a host, a source, and a source type related to the data block. A host field may contain a value identifying a host name or IP address of a device that generated the data. A source field may contain a value identifying a source of the data, such as a pathname of a file or a protocol and port related to received network data. A source type field may contain a value specifying a particular source type label for the data. Additional metadata fields may also be included during the input phase, such as a character encoding of the data, if known, and possibly other values that provide information relevant to later processing steps. In some embodiments, a forwarder forwards the annotated data blocks to another system component (typically an indexer) for further processing.

(110) The data intake and query system allows forwarding of data from one data intake and query instance to another, or even to a third-party system. The data intake and query system can employ different types of forwarders in a configuration.

(111) In some embodiments, a forwarder may contain the essential components needed to forward data. A forwarder can gather data from a variety of inputs and forward the data to an indexer for indexing and searching. A forwarder can also tag metadata (e.g., source, source type, host, etc.).

(112) In some embodiments, a forwarder has the capabilities of the aforementioned forwarder as well as additional capabilities. The forwarder can parse data before forwarding the data (e.g., can associate a time stamp with a portion of data and create an event, etc.) and can route data based on criteria such as source or type of event. The forwarder can also index data locally while forwarding the data to another indexer.

(113) At block **506**, an indexer receives data blocks from a forwarder and parses the data to organize the data into events. In some embodiments, to organize the data into events, an indexer may determine a source type associated with each data block (e.g., by extracting a source type label from the metadata fields associated with the data block, etc.) and refer to a source type configuration corresponding to the identified source type. The source type definition may include one or more properties that indicate to the indexer to automatically determine the boundaries within the received data that indicate the portions of machine data for events. In general, these properties may include regular expression-based rules or delimiter rules where, for example, event boundaries may be indicated by predefined characters or character strings. These predefined characters may include punctuation marks or other special characters including, for example, carriage returns, tabs, spaces, line breaks, etc. If a source type for the data is unknown to the indexer, an indexer may infer a source type for the data by examining the structure of the data. Then, the indexer can apply an inferred source type definition to the data to create the events.

(114) At block **508**, the indexer determines a timestamp for each event. Similar to the process for

parsing machine data, an indexer may again refer to a source type definition associated with the data to locate one or more properties that indicate instructions for determining a timestamp for each event. The properties may, for example, instruct an indexer to extract a time value from a portion of data for the event, to interpolate time values based on timestamps associated with temporally proximate events, to create a timestamp based on a time the portion of machine data was received or generated, to use the timestamp of a previous event, or use any other rules for determining timestamps.

(115) At block **510**, the indexer associates with each event one or more metadata fields including a field containing the timestamp determined for the event. In some embodiments, a timestamp may be included in the metadata fields. These metadata fields may include any number of “default fields” that are associated with all events, and may also include one more custom fields as defined by a user. Similar to the metadata fields associated with the data blocks at block **504**, the default metadata fields associated with each event may include a host, source, and source type field including or in addition to a field storing the timestamp.

(116) At block **512**, an indexer may optionally apply one or more transformations to data included in the events created at block **506**. For example, such transformations can include removing a portion of an event (e.g., a portion used to define event boundaries, extraneous characters from the event, other extraneous text, etc.), masking a portion of an event (e.g., masking a credit card number), removing redundant portions of an event, etc. The transformations applied to events may, for example, be specified in one or more configuration files and referenced by one or more source type definitions.

(117) FIG. 5C illustrates an illustrative example of machine data can be stored in a data store in accordance with various disclosed embodiments. In other embodiments, machine data can be stored in a flat file in a corresponding bucket with an associated index file, such as a time series index or “TSIDX.” As such, the depiction of machine data and associated metadata as rows and columns in the table of FIG. 5C is merely illustrative and is not intended to limit the data format in which the machine data and metadata is stored in various embodiments described herein. In one particular embodiment, machine data can be stored in a compressed or encrypted formatted. In such embodiments, the machine data can be stored with or be associated with data that describes the compression or encryption scheme with which the machine data is stored. The information about the compression or encryption scheme can be used to decompress or decrypt the machine data, and any metadata with which it is stored, at search time.

(118) As mentioned above, certain metadata, e.g., host **536**, source **537**, source type **538** and timestamps **535** can be generated for each event, and associated with a corresponding portion of machine data **539** when storing the event data in a data store, e.g., data store **208**. Any of the metadata can be extracted from the corresponding machine data, or supplied or defined by an entity, such as a user or computer system. The metadata fields can become part of or stored with the event. Note that while the time-stamp metadata field can be extracted from the raw data of each event, the values for the other metadata fields may be determined by the indexer based on information it receives pertaining to the source of the data separate from the machine data.

(119) While certain default or user-defined metadata fields can be extracted from the machine data for indexing purposes, all the machine data within an event can be maintained in its original condition. As such, in embodiments in which the portion of machine data included in an event is unprocessed or otherwise unaltered, it is referred to herein as a portion of raw machine data. In other embodiments, the port of machine data in an event can be processed or otherwise altered. As such, unless certain information needs to be removed for some reasons (e.g., extraneous information, confidential information), all the raw machine data contained in an event can be preserved and saved in its original form. Accordingly, the data store in which the event records are stored is sometimes referred to as a “raw record data store.” The raw record data store contains a record of the raw event data tagged with the various default fields.

(120) In FIG. 5C, the first three rows of the table represent events **531**, **532**, and **533** and are related to a server access log that records requests from multiple clients processed by a server, as indicated by entry of “access.log” in the source column **536**.

(121) In the example shown in FIG. 5C, each of the events **531-533** is associated with a discrete request made from a client device. The raw machine data generated by the server and extracted from a server access log can include the IP address of the client **540**, the user id of the person requesting the document **541**, the time the server finished processing the request **542**, the request line from the client **543**, the status code returned by the server to the client **545**, the size of the object returned to the client (in this case, the gif file requested by the client) **546** and the time spent to serve the request in microseconds **544**. As seen in FIG. 5C, all the raw machine data retrieved from the server access log is retained and stored as part of the corresponding events, **531-533** in the data store.

(122) Event **534** is associated with an entry in a server error log, as indicated by “error.log” in the source column **537** that records errors that the server encountered when processing a client request. Similar to the events related to the server access log, all the raw machine data in the error log file pertaining to event **534** can be preserved and stored as part of the event **534**.

(123) Saving minimally processed or unprocessed machine data in a data store associated with metadata fields in the manner similar to that shown in FIG. 5C is advantageous because it allows search of all the machine data at search time instead of searching only previously specified and identified fields or field-value pairs. As mentioned above, because data structures used by various embodiments of the present disclosure maintain the underlying raw machine data and use a late-binding schema for searching the raw machines data, it enables a user to continue investigating and learn valuable insights about the raw data. In other words, the user is not compelled to know about all the fields of information that will be needed at data ingestion time. As a user learns more about the data in the events, the user can continue to refine the late-binding schema by defining new extraction rules, or modifying or deleting existing extraction rules used by the system.

(124) At blocks **514** and **516**, an indexer can optionally generate a keyword index to facilitate fast keyword searching for events. To build a keyword index, at block **514**, the indexer identifies a set of keywords in each event. At block **516**, the indexer includes the identified keywords in an index, which associates each stored keyword with reference pointers to events containing that keyword (or to locations within events where that keyword is located, other location identifiers, etc.). When an indexer subsequently receives a keyword-based query, the indexer can access the keyword index to quickly identify events containing the keyword.

(125) In some embodiments, the keyword index may include entries for field name-value pairs found in events, where a field name-value pair can include a pair of keywords connected by a symbol, such as an equals sign or colon. This way, events containing these field name-value pairs can be quickly located. In some embodiments, fields can automatically be generated for some or all of the field names of the field name-value pairs at the time of indexing. For example, if the string “dest=10.0.1.2” is found in an event, a field named “dest” may be created for the event, and assigned a value of “10.0.1.2”.

(126) At block **518**, the indexer stores the events with an associated timestamp in a data store **208**. Timestamps enable a user to search for events based on a time range. In some embodiments, the stored events are organized into “buckets,” where each bucket stores events associated with a specific time range based on the timestamps associated with each event. This improves time-based searching, as well as allows for events with recent timestamps, which may have a higher likelihood of being accessed, to be stored in a faster memory to facilitate faster retrieval. For example, buckets containing the most recent events can be stored in flash memory rather than on a hard disk. In some embodiments, each bucket may be associated with an identifier, a time range, and a size constraint.

(127) Each indexer **206** may be responsible for storing and searching a subset of the events contained in a corresponding data store **208**. By distributing events among the indexers and data

stores, the indexers can analyze events for a query in parallel. For example, using mapreduce techniques, each indexer returns partial responses for a subset of events to a search head that combines the results to produce an answer for the query. By storing events in buckets for specific time ranges, an indexer may further optimize the data retrieval process by searching buckets corresponding to time ranges that are relevant to a query. In some embodiments, each bucket may be associated with an identifier, a time range, and a size constraint. In certain embodiments, a bucket can correspond to a file system directory and the machine data, or events, of a bucket can be stored in one or more files of the file system directory. The file system directory can include additional files, such as one or more inverted indexes, high performance indexes, permissions files, configuration files, etc.

(128) In some embodiments, each indexer has a home directory and a cold directory. The home directory of an indexer stores hot buckets and warm buckets, and the cold directory of an indexer stores cold buckets. A hot bucket is a bucket that is capable of receiving and storing events. A warm bucket is a bucket that can no longer receive events for storage but has not yet been moved to the cold directory. A cold bucket is a bucket that can no longer receive events and may be a bucket that was previously stored in the home directory. The home directory may be stored in faster memory, such as flash memory, as events may be actively written to the home directory, and the home directory may typically store events that are more frequently searched and thus are accessed more frequently. The cold directory may be stored in slower and/or larger memory, such as a hard disk, as events are no longer being written to the cold directory, and the cold directory may typically store events that are not as frequently searched and thus are accessed less frequently. In some embodiments, an indexer may also have a quarantine bucket that contains events having potentially inaccurate information, such as an incorrect time stamp associated with the event or a time stamp that appears to be an unreasonable time stamp for the corresponding event. The quarantine bucket may have events from any time range; as such, the quarantine bucket may always be searched at search time. Additionally, an indexer may store old, archived data in a frozen bucket that is not capable of being searched at search time. In some embodiments, a frozen bucket may be stored in slower and/or larger memory, such as a hard disk, and may be stored in offline and/or remote storage.

(129) Moreover, events and buckets can also be replicated across different indexers and data stores to facilitate high availability and disaster recovery as described in U.S. Pat. No. 9,130,971, entitled "SITE-BASED SEARCH AFFINITY", issued on 8 Sep. 2015, and in U.S. patent Ser. No. 14/266,817, entitled "MULTI-SITE CLUSTERING", issued on 1 Sep. 2015, each of which is hereby incorporated by reference in its entirety for all purposes.

(130) FIG. 5B is a block diagram of an example data store **501** that includes a directory for each index (or partition) that contains a portion of data managed by an indexer. FIG. 5B further illustrates details of an embodiment of an inverted index **507B** and an event reference array **515** associated with inverted index **507B**.

(131) The data store **501** can correspond to a data store **208** that stores events managed by an indexer **206** or can correspond to a different data store associated with an indexer **206**. In the illustrated embodiment, the data store **501** includes a main directory **503** associated with a main index and a test directory **505** associated with a test index. However, the data store **501** can include fewer or more directories. In some embodiments, multiple indexes can share a single directory or all indexes can share a common directory. Additionally, although illustrated as a single data store **501**, it will be understood that the data store **501** can be implemented as multiple data stores storing different portions of the information shown in FIG. 5B. For example, a single index or partition can span multiple directories or multiple data stores, and can be indexed or searched by multiple corresponding indexers.

(132) In the illustrated embodiment of FIG. 5B, the index-specific directories **503** and **505** include inverted indexes **507A**, **507B** and **509A**, **509B**, respectively. The inverted indexes **507A** . . . **507B**,

and **509A . . . 509B** can be keyword indexes or field-value pair indexes described herein and can include less or more information that depicted in FIG. 5B.

(133) In some embodiments, the inverted index **507A . . . 507B**, and **509A . . . 509B** can correspond to a distinct time-series bucket that is managed by the indexer **206** and that contains events corresponding to the relevant index (e.g., main index, test index). As such, each inverted index can correspond to a particular range of time for an index. Additional files, such as high performance indexes for each time-series bucket of an index, can also be stored in the same directory as the inverted indexes **507A . . . 507B**, and **509A . . . 509B**. In some embodiments inverted index **507A . . . 507B**, and **509A . . . 509B** can correspond to multiple time-series buckets or inverted indexes **507A . . . 507B**, and **509A . . . 509B** can correspond to a single time-series bucket.

(134) Each inverted index **507A . . . 507B**, and **509A . . . 509B** can include one or more entries, such as keyword (or token) entries or field-value pair entries. Furthermore, in certain embodiments, the inverted indexes **507A . . . 507B**, and **509A . . . 509B** can include additional information, such as a time range **523** associated with the inverted index or an index identifier **525** identifying the index associated with the inverted index **507A . . . 507B**, and **509A . . . 509B**. However, each inverted index **507A . . . 507B**, and **509A . . . 509B** can include less or more information than depicted.

(135) Token entries, such as token entries **511** illustrated in inverted index **507B**, can include a token **511A** (e.g., “error,” “itemID,” etc.) and event references **511B** indicative of events that include the token. For example, for the token “error,” the corresponding token entry includes the token “error” and an event reference, or unique identifier, for each event stored in the corresponding time-series bucket that includes the token “error.” In the illustrated embodiment of FIG. 5B, the error token entry includes the identifiers **3, 5, 6, 8, 11, and 12** corresponding to events managed by the indexer **206** and associated with the index main **503** that are located in the time-series bucket associated with the inverted index **507B**.

(136) In some cases, some token entries can be default entries, automatically determined entries, or user specified entries. In some embodiments, the indexer **206** can identify each word or string in an event as a distinct token and generate a token entry for it. In some cases, the indexer **206** can identify the beginning and ending of tokens based on punctuation, spaces, as described in greater detail herein. In certain cases, the indexer **206** can rely on user input or a configuration file to identify tokens for token entries **511**, etc. It will be understood that any combination of token entries can be included as a default, automatically determined, and/or included based on user-specified criteria.

(137) Similarly, field-value pair entries, such as field-value pair entries **513** shown in inverted index **507B**, can include a field-value pair **513A** and event references **513B** indicative of events that include a field value that corresponds to the field-value pair. For example, for a field-value pair sourcetype:sendmail, a field-value pair entry would include the field-value pair sourcetype:sendmail and a unique identifier, or event reference, for each event stored in the corresponding time-series bucket that includes a sendmail sourcetype.

(138) In some cases, the field-value pair entries **513** can be default entries, automatically determined entries, or user specified entries. As a non-limiting example, the field-value pair entries for the fields host, source, sourcetype can be included in the inverted indexes **507A . . . 507B**, and **509A . . . 509B** as a default. As such, all of the inverted indexes **507A . . . 507B**, and **509A . . . 509B** can include field-value pair entries for the fields host, source, sourcetype. As yet another non-limiting example, the field-value pair entries for the IP\_address field can be user specified and may only appear in the inverted index **507B** based on user-specified criteria. As another non-limiting example, as the indexer indexes the events, it can automatically identify field-value pairs and create field-value pair entries. For example, based on the indexers review of events, it can identify IP\_address as a field in each event and add the IP\_address field-value pair entries to the inverted index **507B**. It will be understood that any combination of field-value pair entries can be included

as a default, automatically determined, or included based on user-specified criteria.

(139) Each unique identifier **517**, or event reference, can correspond to a unique event located in the time series bucket. However, the same event reference can be located in multiple entries. For example, if an event has a sourcetype splunkd, host www1 and token “warning,” then the unique identifier for the event will appear in the field-value pair entries sourcetype:splunkd and host:www1, as well as the token entry “warning.” With reference to the illustrated embodiment of FIG. 5B and the event that corresponds to the event reference **3**, the event reference **3** is found in the field-value pair entries **513** host:hostA, source:sourceB, sourcetype:sourcetypeA, and IP\_address:91.205.189.15 indicating that the event corresponding to the event references is from hostA, sourceB, of sourcetypeA, and includes 91.205.189.15 in the event data.

(140) For some fields, the unique identifier is located in only one field-value pair entry for a particular field. For example, the inverted index may include four sourcetype field-value pair entries corresponding to four different sourcetypes of the events stored in a bucket (e.g., sourcetypes: sendmail, splunkd, web\_access, and web\_service). Within those four sourcetype field-value pair entries, an identifier for a particular event may appear in only one of the field-value pair entries. With continued reference to the example illustrated embodiment of FIG. 5B, since the event reference **7** appears in the field-value pair entry sourcetype:sourcetypeA, then it does not appear in the other field-value pair entries for the sourcetype field, including sourcetype:sourcetypeB, sourcetype:sourcetypeC, and sourcetype:sourcetypeD.

(141) The event references **517** can be used to locate the events in the corresponding bucket. For example, the inverted index can include, or be associated with, an event reference array **515**. The event reference array **515** can include an array entry **517** for each event reference in the inverted index **507B**. Each array entry **517** can include location information **519** of the event corresponding to the unique identifier (non-limiting example: seek address of the event), a timestamp **521** associated with the event, or additional information regarding the event associated with the event reference, etc.

(142) For each token entry **511** or field-value pair entry **513**, the event reference **501B** or unique identifiers can be listed in chronological order or the value of the event reference can be assigned based on chronological data, such as a timestamp associated with the event referenced by the event reference. For example, the event reference **1** in the illustrated embodiment of FIG. 5B can correspond to the first-in-time event for the bucket, and the event reference **12** can correspond to the last-in-time event for the bucket. However, the event references can be listed in any order, such as reverse chronological order, ascending order, descending order, or some other order, etc. Further, the entries can be sorted. For example, the entries can be sorted alphabetically (collectively or within a particular group), by entry origin (e.g., default, automatically generated, user-specified, etc.), by entry type (e.g., field-value pair entry, token entry, etc.), or chronologically by when added to the inverted index, etc. In the illustrated embodiment of FIG. 5B, the entries are sorted first by entry type and then alphabetically.

(143) As a non-limiting example of how the inverted indexes **507A . . . 507B**, and **509A . . . 509B** can be used during a data categorization request command, the indexers can receive filter criteria indicating data that is to be categorized and categorization criteria indicating how the data is to be categorized. Example filter criteria can include, but is not limited to, indexes (or partitions), hosts, sources, sourcetypes, time ranges, field identifier, keywords, etc.

(144) Using the filter criteria, the indexer identifies relevant inverted indexes to be searched. For example, if the filter criteria includes a set of partitions, the indexer can identify the inverted indexes stored in the directory corresponding to the particular partition as relevant inverted indexes. Other means can be used to identify inverted indexes associated with a partition of interest. For example, in some embodiments, the indexer can review an entry in the inverted indexes, such as an index-value pair entry **513** to determine if a particular inverted index is relevant. If the filter criteria does not identify any partition, then the indexer can identify all inverted indexes managed by the

indexer as relevant inverted indexes.

(145) Similarly, if the filter criteria includes a time range, the indexer can identify inverted indexes corresponding to buckets that satisfy at least a portion of the time range as relevant inverted indexes. For example, if the time range is last hour then the indexer can identify all inverted indexes that correspond to buckets storing events associated with timestamps within the last hour as relevant inverted indexes.

(146) When used in combination, an index filter criterion specifying one or more partitions and a time range filter criterion specifying a particular time range can be used to identify a subset of inverted indexes within a particular directory (or otherwise associated with a particular partition) as relevant inverted indexes. As such, the indexer can focus the processing to only a subset of the total number of inverted indexes that the indexer manages.

(147) Once the relevant inverted indexes are identified, the indexer can review them using any additional filter criteria to identify events that satisfy the filter criteria. In some cases, using the known location of the directory in which the relevant inverted indexes are located, the indexer can determine that any events identified using the relevant inverted indexes satisfy an index filter criterion. For example, if the filter criteria includes a partition main, then the indexer can determine that any events identified using inverted indexes within the partition main directory (or otherwise associated with the partition main) satisfy the index filter criterion.

(148) Furthermore, based on the time range associated with each inverted index, the indexer can determine that any events identified using a particular inverted index satisfies a time range filter criterion. For example, if a time range filter criterion is for the last hour and a particular inverted index corresponds to events within a time range of 50 minutes ago to 35 minutes ago, the indexer can determine that any events identified using the particular inverted index satisfy the time range filter criterion. Conversely, if the particular inverted index corresponds to events within a time range of 59 minutes ago to 62 minutes ago, the indexer can determine that some events identified using the particular inverted index may not satisfy the time range filter criterion.

(149) Using the inverted indexes, the indexer can identify event references (and therefore events) that satisfy the filter criteria. For example, if the token “error” is a filter criterion, the indexer can track all event references within the token entry “error.” Similarly, the indexer can identify other event references located in other token entries or field-value pair entries that match the filter criteria. The system can identify event references located in all of the entries identified by the filter criteria. For example, if the filter criteria include the token “error” and field-value pair sourcetype:web\_ui, the indexer can track the event references found in both the token entry “error” and the field-value pair entry sourcetype:web\_ui. As mentioned previously, in some cases, such as when multiple values are identified for a particular filter criterion (e.g., multiple sources for a source filter criterion), the system can identify event references located in at least one of the entries corresponding to the multiple values and in all other entries identified by the filter criteria. The indexer can determine that the events associated with the identified event references satisfy the filter criteria.

(150) In some cases, the indexer can further consult a timestamp associated with the event reference to determine whether an event satisfies the filter criteria. For example, if an inverted index corresponds to a time range that is partially outside of a time range filter criterion, then the indexer can consult a timestamp associated with the event reference to determine whether the corresponding event satisfies the time range criterion. In some embodiments, to identify events that satisfy a time range, the indexer can review an array, such as the event reference array 515 that identifies the time associated with the events. Furthermore, as mentioned above using the known location of the directory in which the relevant inverted indexes are located (or other index identifier), the indexer can determine that any events identified using the relevant inverted indexes satisfy the index filter criterion.

(151) In some cases, based on the filter criteria, the indexer reviews an extraction rule. In certain



embodiments, if the filter criteria includes a field name that does not correspond to a field-value pair entry in an inverted index, the indexer can review an extraction rule, which may be located in a configuration file, to identify a field that corresponds to a field-value pair entry in the inverted index.

(152) For example, the filter criteria includes a field name “sessionID” and the indexer determines that at least one relevant inverted index does not include a field-value pair entry corresponding to the field name sessionID, the indexer can review an extraction rule that identifies how the sessionID field is to be extracted from a particular host, source, or sourcetype (implicitly identifying the particular host, source, or sourcetype that includes a sessionID field). The indexer can replace the field name “sessionID” in the filter criteria with the identified host, source, or sourcetype. In some cases, the field name “sessionID” may be associated with multiples hosts, sources, or sourcetypes, in which case, all identified hosts, sources, and sourcetypes can be added as filter criteria. In some cases, the identified host, source, or sourcetype can replace or be appended to a filter criterion, or be excluded. For example, if the filter criteria includes a criterion for source S1 and the “sessionID” field is found in source S2, the source S2 can replace S1 in the filter criteria, be appended such that the filter criteria includes source S1 and source S2, or be excluded based on the presence of the filter criterion source S1. If the identified host, source, or sourcetype is included in the filter criteria, the indexer can then identify a field-value pair entry in the inverted index that includes a field value corresponding to the identity of the particular host, source, or sourcetype identified using the extraction rule.

(153) Once the events that satisfy the filter criteria are identified, the system, such as the indexer **206** can categorize the results based on the categorization criteria. The categorization criteria can include categories for grouping the results, such as any combination of partition, source, sourcetype, or host, or other categories or fields as desired.

(154) The indexer can use the categorization criteria to identify categorization criteria-value pairs or categorization criteria values by which to categorize or group the results. The categorization criteria-value pairs can correspond to one or more field-value pair entries stored in a relevant inverted index, one or more index-value pairs based on a directory in which the inverted index is located or an entry in the inverted index (or other means by which an inverted index can be associated with a partition), or other criteria-value pair that identifies a general category and a particular value for that category. The categorization criteria values can correspond to the value portion of the categorization criteria-value pair.

(155) As mentioned, in some cases, the categorization criteria-value pairs can correspond to one or more field-value pair entries stored in the relevant inverted indexes. For example, the categorization criteria-value pairs can correspond to field-value pair entries of host, source, and sourcetype (or other field-value pair entry as desired). For instance, if there are ten different hosts, four different sources, and five different sourcetypes for an inverted index, then the inverted index can include ten host field-value pair entries, four source field-value pair entries, and five sourcetype field-value pair entries. The indexer can use the nineteen distinct field-value pair entries as categorization criteria-value pairs to group the results.

(156) Specifically, the indexer can identify the location of the event references associated with the events that satisfy the filter criteria within the field-value pairs, and group the event references based on their location. As such, the indexer can identify the particular field value associated with the event corresponding to the event reference. For example, if the categorization criteria include host and sourcetype, the host field-value pair entries and sourcetype field-value pair entries can be used as categorization criteria-value pairs to identify the specific host and sourcetype associated with the events that satisfy the filter criteria.

(157) In addition, as mentioned, categorization criteria-value pairs can correspond to data other than the field-value pair entries in the relevant inverted indexes. For example, if partition or index is used as a categorization criterion, the inverted indexes may not include partition field-value pair

entries. Rather, the indexer can identify the categorization criteria-value pair associated with the partition based on the directory in which an inverted index is located, information in the inverted index, or other information that associates the inverted index with the partition, etc. As such a variety of methods can be used to identify the categorization criteria-value pairs from the categorization criteria.

(158) Accordingly based on the categorization criteria (and categorization criteria-value pairs), the indexer can generate groupings based on the events that satisfy the filter criteria. As a non-limiting example, if the categorization criteria includes a partition and sourcetype, then the groupings can correspond to events that are associated with each unique combination of partition and sourcetype. For instance, if there are three different partitions and two different sourcetypes associated with the identified events, then the six different groups can be formed, each with a unique partition value-sourcetype value combination. Similarly, if the categorization criteria includes partition, sourcetype, and host and there are two different partitions, three sourcetypes, and five hosts associated with the identified events, then the indexer can generate up to thirty groups for the results that satisfy the filter criteria. Each group can be associated with a unique combination of categorization criteria-value pairs (e.g., unique combinations of partition value sourcetype value, and host value).

(159) In addition, the indexer can count the number of events associated with each group based on the number of events that meet the unique combination of categorization criteria for a particular group (or match the categorization criteria-value pairs for the particular group). With continued reference to the example above, the indexer can count the number of events that meet the unique combination of partition, sourcetype, and host for a particular group.

(160) Each indexer communicates the groupings to the search head. The search head can aggregate the groupings from the indexers and provide the groupings for display. In some cases, the groups are displayed based on at least one of the host, source, sourcetype, or partition associated with the groupings. In some embodiments, the search head can further display the groups based on display criteria, such as a display order or a sort order as described in greater detail above.

(161) As a non-limiting example and with reference to FIG. 5B, consider a request received by an indexer **206** that includes the following filter criteria: keyword=error, partition=main, time range=3/1/17 16:22.00.000-16:28.00.000, sourcetype=sourcetypeC, host=hostB, and the following categorization criteria: source.

(162) Based on the above criteria, the indexer **206** identifies main directory **503** and can ignore\_test directory **505** and any other partition-specific directories. The indexer determines that inverted partition **507B** is a relevant partition based on its location within the main directory **503** and the time range associated with it. For sake of simplicity in this example, the indexer **206** determines that no other inverted indexes in the main directory **503**, such as inverted index **507A** satisfy the time range criterion.

(163) Having identified the relevant inverted index **507B**, the indexer reviews the token entries **511** and the field-value pair entries **513** to identify event references, or events that satisfy all of the filter criteria.

(164) With respect to the token entries **511**, the indexer can review the error token entry and identify event references **3, 5, 6, 8, 11, 12**, indicating that the term “error” is found in the corresponding events. Similarly, the indexer can identify event references **4, 5, 6, 8, 9, 10, 11** in the field-value pair entry sourcetype:sourcetypeC and event references **2, 5, 6, 8, 10, 11** in the field-value pair entry host:hostB. As the filter criteria did not include a source or an IP\_address field-value pair, the indexer can ignore those field-value pair entries.

(165) In addition to identifying event references found in at least one token entry or field-value pair entry (e.g., event references **3, 4, 5, 6, 8, 9, 10, 11, 12**), the indexer can identify events (and corresponding event references) that satisfy the time range criterion using the event reference array **1614** (e.g., event references **2, 3, 4, 5, 6, 7, 8, 9, 10**). Using the information obtained from the

inverted index **507B** (including the event reference array **515**), the indexer **206** can identify the event references that satisfy all of the filter criteria (e.g., event references **5, 6, 8**).

(166) Having identified the events (and event references) that satisfy all of the filter criteria, the indexer **206** can group the event references using the received categorization criteria (source). In doing so, the indexer can determine that event references **5** and **6** are located in the field-value pair entry source:sourceD (or have matching categorization criteria-value pairs) and event reference **8** is located in the field-value pair entry source:sourceC. Accordingly, the indexer can generate a sourceC group having a count of one corresponding to reference **8** and a sourceD group having a count of two corresponding to references **5** and **6**. This information can be communicated to the search head. In turn the search head can aggregate the results from the various indexers and display the groupings. As mentioned above, in some embodiments, the groupings can be displayed based at least in part on the categorization criteria, including at least one of host, source, sourcetype, or partition.

(167) It will be understood that a change to any of the filter criteria or categorization criteria can result in different groupings. As a one non-limiting example, a request received by an indexer **206** that includes the following filter criteria: partition-main, time range=3/1/17 3/1/17 16:21:20.000-16:28:17.000, and the following categorization criteria: host, source, sourcetype would result in the indexer identifying event references **1-12** as satisfying the filter criteria. The indexer would then generate up to 24 groupings corresponding to the 24 different combinations of the categorization criteria-value pairs, including host (hostA, hostB), source (sourceA, sourceB, sourceC, sourceD), and sourcetype (sourcetypeA, sourcetypeB, sourcetypeC). However, as there are only twelve events identifiers in the illustrated embodiment and some fall into the same grouping, the indexer generates eight groups and counts as follows: Group 1 (hostA, sourceA, sourcetypeA): 1 (event reference **7**) Group 2 (hostA, sourceA, sourcetypeB): 2 (event references **1, 12**) Group 3 (hostA, sourceA, sourcetypeC): 1 (event reference **4**) Group 4 (hostA, sourceB, sourcetypeA): 1 (event reference **3**) Group 5 (hostA, sourceB, sourcetypeC): 1 (event reference **9**) Group 6 (hostB, sourceC, sourcetypeA): 1 (event reference **2**) Group 7 (hostB, sourceC, sourcetypeC): 2 (event references **8, 11**) Group 8 (hostB, sourceD, sourcetypeC): 3 (event references **5, 6, 10**)

(168) As noted, each group has a unique combination of categorization criteria-value pairs or categorization criteria values. The indexer communicates the groups to the search head for aggregation with results received from other indexers. In communicating the groups to the search head, the indexer can include the categorization criteria-value pairs for each group and the count. In some embodiments, the indexer can include more or less information. For example, the indexer can include the event references associated with each group and other identifying information, such as the indexer or inverted index used to identify the groups.

(169) As another non-limiting examples, a request received by an indexer **206** that includes the following filter criteria: partition=main, time range=3/1/17 3/1/17 16:21:20.000-16:28:17.000, source=sourceA, sourceD, and keyword-itemID and the following categorization criteria: host, source, sourcetype would result in the indexer identifying event references **4, 7, and 10** as satisfying the filter criteria, and generate the following groups: Group 1 (hostA, sourceA, sourcetypeC): 1 (event reference **4**) Group 2 (hostA, sourceA, sourcetypeA): 1 (event reference **7**) Group 3 (hostB, sourceD, sourcetypeC): 1 (event references **10**)

(170) The indexer communicates the groups to the search head for aggregation with results received from other indexers. As will be understand there are myriad ways for filtering and categorizing the events and event references. For example, the indexer can review multiple inverted indexes associated with a partition or review the inverted indexes of multiple partitions, and categorize the data using any one or any combination of partition, host, source, sourcetype, or other category, as desired.

(171) Further, if a user interacts with a particular group, the indexer can provide additional information regarding the group. For example, the indexer can perform a targeted search or

sampling of the events that satisfy the filter criteria and the categorization criteria for the selected group, also referred to as the filter criteria corresponding to the group or filter criteria associated with the group.

(172) In some cases, to provide the additional information, the indexer relies on the inverted index. For example, the indexer can identify the event references associated with the events that satisfy the filter criteria and the categorization criteria for the selected group and then use the event reference array **515** to access some or all of the identified events. In some cases, the categorization criteria values or categorization criteria-value pairs associated with the group become part of the filter criteria for the review.

(173) With reference to FIG. **5B** for instance, suppose a group is displayed with a count of six corresponding to event references **4, 5, 6, 8, 10, 11** (i.e., event references **4, 5, 6, 8, 10, 11** satisfy the filter criteria and are associated with matching categorization criteria values or categorization criteria-value pairs) and a user interacts with the group (e.g., selecting the group, clicking on the group, etc.). In response, the search head communicates with the indexer to provide additional information regarding the group.

(174) In some embodiments, the indexer identifies the event references associated with the group using the filter criteria and the categorization criteria for the group (e.g., categorization criteria values or categorization criteria-value pairs unique to the group). Together, the filter criteria and the categorization criteria for the group can be referred to as the filter criteria associated with the group. Using the filter criteria associated with the group, the indexer identifies event references **4, 5, 6, 8, 10, 11**.

(175) Based on a sampling criteria, discussed in greater detail above, the indexer can determine that it will analyze a sample of the events associated with the event references **4, 5, 6, 8, 10, 11**. For example, the sample can include analyzing event data associated with the event references **5, 8, 10**. In some embodiments, the indexer can use the event reference array **515** to access the event data associated with the event references **5, 8, 10**. Once accessed, the indexer can compile the relevant information and provide it to the search head for aggregation with results from other indexers. By identifying events and sampling event data using the inverted indexes, the indexer can reduce the amount of actual data this is analyzed and the number of events that are accessed in order to generate the summary of the group and provide a response in less time.

(176) FIG. **6A** is a flow diagram of an example method that illustrates how a search head and indexers perform a search query, in accordance with example embodiments. At block **602**, a search head receives a search query from a client. At block **604**, the search head analyzes the search query to determine what portion(s) of the query can be delegated to indexers and what portions of the query can be executed locally by the search head. At block **606**, the search head distributes the determined portions of the query to the appropriate indexers. In some embodiments, a search head cluster may take the place of an independent search head where each search head in the search head cluster coordinates with peer search heads in the search head cluster to schedule jobs, replicate search results, update configurations, fulfill search requests, etc. In some embodiments, the search head (or each search head) communicates with a master node (also known as a cluster master, not shown in FIG. **2**) that provides the search head with a list of indexers to which the search head can distribute the determined portions of the query. The master node maintains a list of active indexers and can also designate which indexers may have responsibility for responding to queries over certain sets of events. A search head may communicate with the master node before the search head distributes queries to indexers to discover the addresses of active indexers.

(177) At block **608**, the indexers to which the query was distributed, search data stores associated with them for events that are responsive to the query. To determine which events are responsive to the query, the indexer searches for events that match the criteria specified in the query. These criteria can include matching keywords or specific values for certain fields. The searching operations at block **608** may use the late-binding schema to extract values for specified fields from

events at the time the query is processed. In some embodiments, one or more rules for extracting field values may be specified as part of a source type definition in a configuration file. The indexers may then either send the relevant events back to the search head, or use the events to determine a partial result, and send the partial result back to the search head.

(178) At block **610**, the search head combines the partial results and/or events received from the indexers to produce a final result for the query. In some examples, the results of the query are indicative of performance or security of the IT environment and may help improve the performance of components in the IT environment. This final result may comprise different types of data depending on what the query requested. For example, the results can include a listing of matching events returned by the query, or some type of visualization of the data from the returned events. In another example, the final result can include one or more calculated values derived from the matching events.

(179) The results generated by the system **108** can be returned to a client using different techniques. For example, one technique streams results or relevant events back to a client in real-time as they are identified. Another technique waits to report the results to the client until a complete set of results (which may include a set of relevant events or a result based on relevant events) is ready to return to the client. Yet another technique streams interim results or relevant events back to the client in real-time until a complete set of results is ready, and then returns the complete set of results to the client. In another technique, certain results are stored as “search jobs” and the client may retrieve the results by referring the search jobs.

(180) The search head can also perform various operations to make the search more efficient. For example, before the search head begins execution of a query, the search head can determine a time range for the query and a set of common keywords that all matching events include. The search head may then use these parameters to query the indexers to obtain a superset of the eventual results. Then, during a filtering stage, the search head can perform field-extraction operations on the superset to produce a reduced set of search results. This speeds up queries, which may be particularly helpful for queries that are performed on a periodic basis.

(181) Various embodiments of the present disclosure can be implemented using, or in conjunction with, a pipelined command language. A pipelined command language is a language in which a set of inputs or data is operated on by a first command in a sequence of commands, and then subsequent commands in the order they are arranged in the sequence. Such commands can include any type of functionality for operating on data, such as retrieving, searching, filtering, aggregating, processing, transmitting, and the like. As described herein, a query can thus be formulated in a pipelined command language and include any number of ordered or unordered commands for operating on data.

(182) Splunk Processing Language (SPL) is an example of a pipelined command language in which a set of inputs or data is operated on by any number of commands in a particular sequence. A sequence of commands, or command sequence, can be formulated such that the order in which the commands are arranged defines the order in which the commands are applied to a set of data or the results of an earlier executed command. For example, a first command in a command sequence can operate to search or filter for specific data in particular set of data. The results of the first command can then be passed to another command listed later in the command sequence for further processing.

(183) In various embodiments, a query can be formulated as a command sequence defined in a command line of a search UI. In some embodiments, a query can be formulated as a sequence of SPL commands. Some or all of the SPL commands in the sequence of SPL commands can be separated from one another by a pipe symbol “|”. In such embodiments, a set of data, such as a set of events, can be operated on by a first SPL command in the sequence, and then a subsequent SPL command following a pipe symbol “|” after the first SPL command operates on the results produced by the first SPL command or other set of data, and so on for any additional SPL

commands in the sequence. As such, a query formulated using SPL comprises a series of consecutive commands that are delimited by pipe “|” characters. The pipe character indicates to the system that the output or result of one command (to the left of the pipe) should be used as the input for one of the subsequent commands (to the right of the pipe). This enables formulation of queries defined by a pipeline of sequenced commands that refines or enhances the data at each step along the pipeline until the desired results are attained. Accordingly, various embodiments described herein can be implemented with Splunk Processing Language (SPL) used in conjunction with the SPLUNK® ENTERPRISE system.

(184) While a query can be formulated in many ways, a query can start with a search command and one or more corresponding search terms at the beginning of the pipeline. Such search terms can include any combination of keywords, phrases, times, dates, Boolean expressions, fieldname-field value pairs, etc. that specify which results should be obtained from an index. The results can then be passed as inputs into subsequent commands in a sequence of commands by using, for example, a pipe character. The subsequent commands in a sequence can include directives for additional processing of the results once it has been obtained from one or more indexes. For example, commands may be used to filter unwanted information out of the results, extract more information, evaluate field values, calculate statistics, reorder the results, create an alert, create summary of the results, or perform some type of aggregation function. In some embodiments, the summary can include a graph, chart, metric, or other visualization of the data. An aggregation function can include analysis or calculations to return an aggregate value, such as an average value, a sum, a maximum value, a root mean square, statistical values, and the like.

(185) Due to its flexible nature, use of a pipelined command language in various embodiments is advantageous because it can perform “filtering” as well as “processing” functions. In other words, a single query can include a search command and search term expressions, as well as data-analysis expressions. For example, a command at the beginning of a query can perform a “filtering” step by retrieving a set of data based on a condition (e.g., records associated with server response times of less than 1 microsecond). The results of the filtering step can then be passed to a subsequent command in the pipeline that performs a “processing” step (e.g., calculation of an aggregate value related to the filtered events such as the average response time of servers with response times of less than 1 microsecond). Furthermore, the search command can allow events to be filtered by keyword as well as field value criteria. For example, a search command can filter out all events containing the word “warning” or filter out all events where a field value associated with a field “clientip” is “10.0.1.2.”

(186) The results obtained or generated in response to a command in a query can be considered a set of results data. The set of results data can be passed from one command to another in any data format. In one embodiment, the set of result data can be in the form of a dynamically created table. Each command in a particular query can redefine the shape of the table. In some implementations, an event retrieved from an index in response to a query can be considered a row with a column for each field value. Columns contain basic information about the data and also may contain data that has been dynamically extracted at search time.

(187) FIG. 6B provides a visual representation of the manner in which a pipelined command language or query operates in accordance with the disclosed embodiments. The query **630** can be inputted by the user into a search. The query comprises a search, the results of which are piped to two commands (namely, command **1** and command **2**) that follow the search step.

(188) Disk **622** represents the event data in the raw record data store.

(189) When a user query is processed, a search step will precede other queries in the pipeline in order to generate a set of events at block **640**. For example, the query can comprise search terms “sourcetype=syslog ERROR” at the front of the pipeline as shown in FIG. 6B. Intermediate results table **624** shows fewer rows because it represents the subset of events retrieved from the index that matched the search terms “sourcetype=syslog ERROR” from search command **630**. By way of

further example, instead of a search step, the set of events at the head of the pipeline may be generating by a call to a pre-existing inverted index (as will be explained later).

(190) At block **642**, the set of events generated in the first part of the query may be piped to a query that searches the set of events for field-value pairs or for keywords. For example, the second intermediate results table **626** shows fewer columns, representing the result of the top command, “top user” which summarizes the events into a list of the top 10 users and displays the user, count, and percentage.

(191) Finally, at block **644**, the results of the prior stage can be pipelined to another stage where further filtering or processing of the data can be performed, e.g., preparing the data for display purposes, filtering the data based on a condition, performing a mathematical calculation with the data, etc. As shown in FIG. **6B**, the “fields-percent” part of command **630** removes the column that shows the percentage, thereby, leaving a final results table **628** without a percentage column. In different embodiments, other query languages, such as the Structured Query Language (“SQL”), can be used to create a query.

(192) The search head **210** allows users to search and visualize events generated from machine data received from homogenous data sources. The search head **210** also allows users to search and visualize events generated from machine data received from heterogeneous data sources. The search head **210** includes various mechanisms, which may additionally reside in an indexer **206**, for processing a query. A query language may be used to create a query, such as any suitable pipelined query language. For example, Splunk Processing Language (SPL) can be utilized to make a query. SPL is a pipelined search language in which a set of inputs is operated on by a first command in a command line, and then a subsequent command following the pipe symbol “|” operates on the results produced by the first command, and so on for additional commands. Other query languages, such as the Structured Query Language (“SQL”), can be used to create a query.

(193) In response to receiving the search query, search head **210** uses extraction rules to extract values for fields in the events being searched. The search head **210** obtains extraction rules that specify how to extract a value for fields from an event. Extraction rules can comprise regex rules that specify how to extract values for the fields corresponding to the extraction rules. In addition to specifying how to extract field values, the extraction rules may also include instructions for deriving a field value by performing a function on a character string or value retrieved by the extraction rule. For example, an extraction rule may truncate a character string or convert the character string into a different data format. In some cases, the query itself can specify one or more extraction rules.

(194) The search head **210** can apply the extraction rules to events that it receives from indexers **206**. Indexers **206** may apply the extraction rules to events in an associated data store **208**. Extraction rules can be applied to all the events in a data store or to a subset of the events that have been filtered based on some criteria (e.g., event time stamp values, etc.). Extraction rules can be used to extract one or more values for a field from events by parsing the portions of machine data in the events and examining the data for one or more patterns of characters, numbers, delimiters, etc., that indicate where the field begins and, optionally, ends.

(195) FIG. **7A** is a diagram of an example scenario where a common customer identifier is found among log data received from three disparate data sources, in accordance with example embodiments. In this example, a user submits an order for merchandise using a vendor's shopping application program **701** running on the user's system. In this example, the order was not delivered to the vendor's server due to a resource exception at the destination server that is detected by the middleware code **702**. The user then sends a message to the customer support server **703** to complain about the order failing to complete. The three systems **701**, **702**, and **703** are disparate systems that do not have a common logging format. The order application **701** sends log data **704** to the data intake and query system in one format, the middleware code **702** sends error log data **705** in a second format, and the support server **703** sends log data **706** in a third format.

(196) Using the log data received at one or more indexers **206** from the three systems, the vendor can uniquely obtain an insight into user activity, user experience, and system behavior. The search head **210** allows the vendor's administrator to search the log data from the three systems that one or more indexers **206** are responsible for searching, thereby obtaining correlated information, such as the order number and corresponding customer ID number of the person placing the order. The system also allows the administrator to see a visualization of related events via a user interface. The administrator can query the search head **210** for customer ID field value matches across the log data from the three systems that are stored at the one or more indexers **206**. The customer ID field value exists in the data gathered from the three systems, but the customer ID field value may be located in different areas of the data given differences in the architecture of the systems. There is a semantic relationship between the customer ID field values generated by the three systems. The search head **210** requests events from the one or more indexers **206** to gather relevant events from the three systems. The search head **210** then applies extraction rules to the events in order to extract field values that it can correlate. The search head may apply a different extraction rule to each set of events from each system when the event format differs among systems. In this example, the user interface can display to the administrator the events corresponding to the common customer ID field values **707**, **708**, and **709**, thereby providing the administrator with insight into a customer's experience.

(197) Note that query results can be returned to a client, a search head, or any other system component for further processing. In general, query results may include a set of one or more events, a set of one or more values obtained from the events, a subset of the values, statistics calculated based on the values, a report containing the values, a visualization (e.g., a graph or chart) generated from the values, and the like.

(198) The search system enables users to run queries against the stored data to retrieve events that meet criteria specified in a query, such as containing certain keywords or having specific values in defined fields. FIG. 7B illustrates the manner in which keyword searches and field searches are processed in accordance with disclosed embodiments.

(199) If a user inputs a search query into search bar **710** that includes only keywords (also known as “tokens”), e.g., the keyword “error” or “warning”, the query search engine of the data intake and query system searches for those keywords directly in the event data **711** stored in the raw record data store. Note that while FIG. 7B only illustrates four events **712**, **713**, **714**, **715**, the raw record data store (corresponding to data store **208** in FIG. 2) may contain records for millions of events.

(200) As disclosed above, an indexer can optionally generate a keyword index to facilitate fast keyword searching for event data. The indexer includes the identified keywords in an index, which associates each stored keyword with reference pointers to events containing that keyword (or to locations within events where that keyword is located, other location identifiers, etc.). When an indexer subsequently receives a keyword-based query, the indexer can access the keyword index to quickly identify events containing the keyword. For example, if the keyword “HTTP” was indexed by the indexer at index time, and the user searches for the keyword “HTTP,” the events **712**, **713**, and **714**, will be identified based on the results returned from the keyword index. As noted above, the index contains reference pointers to the events containing the keyword, which allows for efficient retrieval of the relevant events from the raw record data store.

(201) If a user searches for a keyword that has not been indexed by the indexer, the data intake and query system would nevertheless be able to retrieve the events by searching the event data for the keyword in the raw record data store directly as shown in FIG. 7B. For example, if a user searches for the keyword “frank”, and the name “frank” has not been indexed at index time, the data intake and query system will search the event data directly and return the first event **712**. Note that whether the keyword has been indexed at index time or not, in both cases the raw data of the events **712-715** is accessed from the raw data record store to service the keyword search. In the case where the keyword has been indexed, the index will contain a reference pointer that will allow for a more



efficient retrieval of the event data from the data store. If the keyword has not been indexed, the search engine will need to search through all the records in the data store to service the search. (202) In most cases, however, in addition to keywords, a user's search will also include fields. The term "field" refers to a location in the event data containing one or more values for a specific data item. Often, a field is a value with a fixed, delimited position on a line, or a name and value pair, where there is a single value to each field name. A field can also be multivalued, that is, it can appear more than once in an event and have a different value for each appearance, e.g., email address fields. Fields are searchable by the field name or field name-value pairs. Some examples of fields are "clientip" for IP addresses accessing a web server, or the "From" and "To" fields in email addresses.

(203) By way of further example, consider the search, "status=404". This search query finds events with "status" fields that have a value of "404." When the search is run, the search engine does not look for events with any other "status" value. It also does not look for events containing other fields that share "404" as a value. As a result, the search returns a set of results that are more focused than if "404" had been used in the search string as part of a keyword search. Note also that fields can appear in events as "key=value" pairs such as "user\_name=Bob." But in most cases, field values appear in fixed, delimited positions without identifying keys. For example, the data store may contain events where the "user\_name" value always appears by itself after the timestamp as illustrated by the following string: "Nov 15 09:33:22 johnmedlock."

(204) The data intake and query system advantageously allows for search time field extraction. In other words, fields can be extracted from the event data at search time using late-binding schema as opposed to at data ingestion time, which was a major limitation of the prior art systems.

(205) In response to receiving the search query, search head **210** uses extraction rules to extract values for the fields associated with a field or fields in the event data being searched. The search head **210** obtains extraction rules that specify how to extract a value for certain fields from an event. Extraction rules can comprise regex rules that specify how to extract values for the relevant fields. In addition to specifying how to extract field values, the extraction rules may also include instructions for deriving a field value by performing a function on a character string or value retrieved by the extraction rule. For example, a transformation rule may truncate a character string, or convert the character string into a different data format. In some cases, the query itself can specify one or more extraction rules.

(206) FIG. 7B illustrates the manner in which configuration files may be used to configure custom fields at search time in accordance with the disclosed embodiments. In response to receiving a search query, the data intake and query system determines if the query references a "field." For example, a query may request a list of events where the "clientip" field equals "127.0.0.1." If the query itself does not specify an extraction rule and if the field is not a metadata field, e.g., time, host, source, source type, etc., then in order to determine an extraction rule, the search engine may, in one or more embodiments, need to locate configuration file **716** during the execution of the search as shown in FIG. 7B.

(207) Configuration file **716** may contain extraction rules for all the various fields that are not metadata fields, e.g., the "clientip" field. The extraction rules may be inserted into the configuration file in a variety of ways. In some embodiments, the extraction rules can comprise regular expression rules that are manually entered in by the user. Regular expressions match patterns of characters in text and are used for extracting custom fields in text.

(208) In one or more embodiments, as noted above, a field extractor may be configured to automatically generate extraction rules for certain field values in the events when the events are being created, indexed, or stored, or possibly at a later time. In one embodiment, a user may be able to dynamically create custom fields by highlighting portions of a sample event that should be extracted as fields using a graphical user interface. The system would then generate a regular expression that extracts those fields from similar events and store the regular expression as an

extraction rule for the associated field in the configuration file **716**.

(209) In some embodiments, the indexers may automatically discover certain custom fields at index time and the regular expressions for those fields will be automatically generated at index time and stored as part of extraction rules in configuration file **716**. For example, fields that appear in the event data as “key=value” pairs may be automatically extracted as part of an automatic field discovery process. Note that there may be several other ways of adding field definitions to configuration files in addition to the methods discussed herein.

(210) The search head **210** can apply the extraction rules derived from configuration file **716** to event data that it receives from indexers **206**. Indexers **206** may apply the extraction rules from the configuration file to events in an associated data store **208**. Extraction rules can be applied to all the events in a data store, or to a subset of the events that have been filtered based on some criteria (e.g., event time stamp values, etc.). Extraction rules can be used to extract one or more values for a field from events by parsing the event data and examining the event data for one or more patterns of characters, numbers, delimiters, etc., that indicate where the field begins and, optionally, ends.

(211) In one more embodiments, the extraction rule in configuration file **716** will also need to define the type or set of events that the rule applies to. Because the raw record data store will contain events from multiple heterogeneous sources, multiple events may contain the same fields in different locations because of discrepancies in the format of the data generated by the various sources. Furthermore, certain events may not contain a particular field at all. For example, event **715** also contains “clientip” field, however, the “clientip” field is in a different format from the events **712**, **713**, and **714**. To address the discrepancies in the format and content of the different types of events, the configuration file will also need to specify the set of events that an extraction rule applies to, e.g., extraction rule **717** specifies a rule for filtering by the type of event and contains a regular expression for parsing out the field value. Accordingly, each extraction rule will pertain to only a particular type of event. If a particular field, e.g., “clientip” occurs in multiple events, each of those types of events would need its own corresponding extraction rule in the configuration file **716** and each of the extraction rules would comprise a different regular expression to parse out the associated field value. The most common way to categorize events is by source type because events generated by a particular source can have the same format.

(212) The field extraction rules stored in configuration file **716** perform search-time field extractions. For example, for a query that requests a list of events with source type “access\_combined” where the “clientip” field equals “127.0.0.1,” the query search engine would first locate the configuration file **716** to retrieve extraction rule **717** that would allow it to extract values associated with the “clientip” field from the event data **720** “where the source type is “access\_combined. After the “clientip” field has been extracted from all the events comprising the “clientip” field where the source type is “access\_combined,” the query search engine can then execute the field criteria by performing the compare operation to filter out the events where the “clientip” field equals “127.0.0.1.” In the example shown in FIG. 7B, the events **712**, **713**, and **714** would be returned in response to the user query. In this manner, the search engine can service queries containing field criteria in addition to queries containing keyword criteria (as explained above).

(213) The configuration file can be created during indexing. It may either be manually created by the user or automatically generated with certain predetermined field extraction rules. As discussed above, the events may be distributed across several indexers, wherein each indexer may be responsible for storing and searching a subset of the events contained in a corresponding data store. In a distributed indexer system, each indexer would need to maintain a local copy of the configuration file that is synchronized periodically across the various indexers.

(214) The ability to add schema to the configuration file at search time results in increased efficiency. A user can create new fields at search time and simply add field definitions to the configuration file. As a user learns more about the data in the events, the user can continue to refine

the late-binding schema by adding new fields, deleting fields, or modifying the field extraction rules in the configuration file for use the next time the schema is used by the system. Because the data intake and query system maintains the underlying raw data and uses late-binding schema for searching the raw data, it enables a user to continue investigating and learn valuable insights about the raw data long after data ingestion time.

(215) The ability to add multiple field definitions to the configuration file at search time also results in increased flexibility. For example, multiple field definitions can be added to the configuration file to capture the same field across events generated by different source types. This allows the data intake and query system to search and correlate data across heterogeneous sources flexibly and efficiently.

(216) Further, by providing the field definitions for the queried fields at search time, the configuration file **716** allows the record data store to be field searchable. In other words, the raw record data store can be searched using keywords as well as fields, wherein the fields are searchable name/value pairings that distinguish one event from another and can be defined in configuration file **716** using extraction rules. In comparison to a search containing field names, a keyword search does not need the configuration file and can search the event data directly as shown in FIG. **7B**.

(217) It should also be noted that any events filtered out by performing a search-time field extraction using a configuration file can be further processed by directing the results of the filtering step to a processing step using a pipelined search language. Using the prior example, a user could pipeline the results of the compare step to an aggregate function by asking the query search engine to count the number of events where the “clientip” field equals “127.0.0.1.”

(218) FIG. **8A** is an interface diagram of an example user interface for a search screen **800**, in accordance with example embodiments. Search screen **800** includes a search bar **802** that accepts user input in the form of a search string. It also includes a time range picker **812** that enables the user to specify a time range for the search. For historical searches (e.g., searches based on a particular historical time range), the user can select a specific time range, or alternatively a relative time range, such as “today,” “yesterday” or “last week.” For real-time searches (e.g., searches whose results are based on data received in real-time), the user can select the size of a preceding time window to search for real-time events. Search screen **800** also initially displays a “data summary” dialog as is illustrated in FIG. **8B** that enables the user to select different sources for the events, such as by selecting specific hosts and log files.

(219) After the search is executed, the search screen **800** in FIG. **8A** can display the results through search results tabs **804**, wherein search results tabs **804** includes: an “events tab” that displays various information about events returned by the search; a “statistics tab” that displays statistics about the search results; and a “visualization tab” that displays various visualizations of the search results. The events tab illustrated in FIG. **8A** displays a timeline graph **805** that graphically illustrates the number of events that occurred in one-hour intervals over the selected time range. The events tab also displays an events list **808** that enables a user to view the machine data in each of the returned events.

(220) The events tab additionally displays a sidebar that is an interactive field picker **806**. The field picker **806** may be displayed to a user in response to the search being executed and allows the user to further analyze the search results based on the fields in the events of the search results. The field picker **806** includes field names that reference fields present in the events in the search results. The field picker may display any Selected Fields **820** that a user has pre-selected for display (e.g., host, source, sourcetype) and may also display any Interesting Fields **822** that the system determines may be interesting to the user based on pre-specified criteria (e.g., action, bytes, categoryid, clientip, date\_hour, date\_mday, date\_minute, etc.). The field picker also provides an option to display field names for all the fields present in the events of the search results using the All Fields control **824**.

(221) Each field name in the field picker **806** has a value type identifier to the left of the field name, such as value type identifier **826**. A value type identifier identifies the type of value for the respective field, such as an “a” for fields that include literal values or a “#” for fields that include numerical values.

(222) Each field name in the field picker also has a unique value count to the right of the field name, such as unique value count **828**. The unique value count indicates the number of unique values for the respective field in the events of the search results.

(223) Each field name is selectable to view the events in the search results that have the field referenced by that field name. For example, a user can select the “host” field name, and the events shown in the events list **808** will be updated with events in the search results that have the field that is reference by the field name “host.”

(224) A data model is a hierarchically structured search-time mapping of semantic knowledge about one or more datasets. It encodes the domain knowledge used to build a variety of specialized searches of those datasets. Those searches, in turn, can be used to generate reports.

(225) A data model is composed of one or more “objects” (or “data model objects”) that define or otherwise correspond to a specific set of data. An object is defined by constraints and attributes. An object's constraints are search criteria that define the set of events to be operated on by running a search having that search criteria at the time the data model is selected. An object's attributes are the set of fields to be exposed for operating on that set of events generated by the search criteria.

(226) Objects in data models can be arranged hierarchically in parent/child relationships. Each child object represents a subset of the dataset covered by its parent object. The top-level objects in data models are collectively referred to as “root objects.”

(227) Child objects have inheritance. Child objects inherit constraints and attributes from their parent objects and may have additional constraints and attributes of their own. Child objects provide a way of filtering events from parent objects. Because a child object may provide an additional constraint in addition to the constraints it has inherited from its parent object, the dataset it represents may be a subset of the dataset that its parent represents. For example, a first data model object may define a broad set of data pertaining to e-mail activity generally, and another data model object may define specific datasets within the broad dataset, such as a subset of the e-mail data pertaining specifically to e-mails sent. For example, a user can simply select an “e-mail activity” data model object to access a dataset relating to e-mails generally (e.g., sent or received), or select an “e-mails sent” data model object (or data sub-model object) to access a dataset relating to e-mails sent.

(228) Because a data model object is defined by its constraints (e.g., a set of search criteria) and attributes (e.g., a set of fields), a data model object can be used to quickly search data to identify a set of events and to identify a set of fields to be associated with the set of events. For example, an “e-mails sent” data model object may specify a search for events relating to e-mails that have been sent, and specify a set of fields that are associated with the events. Thus, a user can retrieve and use the “e-mails sent” data model object to quickly search source data for events relating to sent e-mails, and may be provided with a listing of the set of fields relevant to the events in a user interface screen.

(229) Examples of data models can include electronic mail, authentication, databases, intrusion detection, malware, application state, alerts, compute inventory, network sessions, network traffic, performance, audits, updates, vulnerabilities, etc. Data models and their objects can be designed by knowledge managers in an organization, and they can enable downstream users to quickly focus on a specific set of data. A user iteratively applies a model development tool (not shown in FIG. **8A**) to prepare a query that defines a subset of events and assigns an object name to that subset. A child subset is created by further limiting a query that generated a parent subset.

(230) Data definitions in associated schemas can be taken from the common information model (CIM) or can be devised for a particular schema and optionally added to the CIM. Child objects

inherit fields from parents and can include fields not present in parents. A model developer can select fewer extraction rules than are available for the sources returned by the query that defines events belonging to a model. Selecting a limited set of extraction rules can be a tool for simplifying and focusing the data model, while allowing a user flexibility to explore the data subset. Development of a data model is further explained in U.S. Pat. Nos. 8,788,525 and 8,788,526, both entitled “DATA MODEL FOR MACHINE DATA FOR SEMANTIC SEARCH”, both issued on 22 Jul. 2014, U.S. Pat. No. 8,983,994, entitled “GENERATION OF A DATA MODEL FOR SEARCHING MACHINE DATA”, issued on 17 Mar. 2015, U.S. Pat. No. 9,128,980, entitled “GENERATION OF A DATA MODEL APPLIED TO QUERIES”, issued on 8 Sep. 2015, and U.S. Pat. No. 9,589,012, entitled “GENERATION OF A DATA MODEL APPLIED TO OBJECT QUERIES”, issued on 7 Mar. 2017, each of which is hereby incorporated by reference in its entirety for all purposes.

(231) A data model can also include reports. One or more report formats can be associated with a particular data model and be made available to run against the data model. A user can use child objects to design reports with object datasets that already have extraneous data pre-filtered out. In some embodiments, the data intake and query system **108** provides the user with the ability to produce reports (e.g., a table, chart, visualization, etc.) without having to enter SPL, SQL, or other query language terms into a search screen. Data models are used as the basis for the search feature.

(232) Data models may be selected in a report generation interface. The report generator supports drag-and-drop organization of fields to be summarized in a report. When a model is selected, the fields with available extraction rules are made available for use in the report. The user may refine and/or filter search results to produce more precise reports. The user may select some fields for organizing the report and select other fields for providing detail according to the report organization. For example, “region” and “salesperson” are fields used for organizing the report and sales data can be summarized (subtotaled and totaled) within this organization. The report generator allows the user to specify one or more fields within events and apply statistical analysis on values extracted from the specified one or more fields. The report generator may aggregate search results across sets of events and generate statistics based on aggregated search results. Building reports using the report generation interface is further explained in U.S. patent application Ser. No. 14/503,335, entitled “GENERATING REPORTS FROM UNSTRUCTURED DATA,” filed on 30 Sep. 2014, and which is hereby incorporated by reference in its entirety for all purposes. Data visualizations also can be generated in a variety of formats, by reference to the data model. Reports, data visualizations, and data model objects can be saved and associated with the data model for future use. The data model object may be used to perform searches of other data.

(233) FIGS. **9-15** are interface diagrams of example report generation user interfaces, in accordance with example embodiments. The report generation process may be driven by a predefined data model object, such as a data model object defined and/or saved via a reporting application or a data model object obtained from another source. A user can load a saved data model object using a report editor. For example, the initial search query and fields used to drive the report editor may be obtained from a data model object. The data model object that is used to drive a report generation process may define a search and a set of fields. Upon loading of the data model object, the report generation process may enable a user to use the fields (e.g., the fields defined by the data model object) to define criteria for a report (e.g., filters, split rows/columns, aggregates, etc.) and the search may be used to identify events (e.g., to identify events responsive to the search) used to generate the report. That is, for example, if a data model object is selected to drive a report editor, the graphical user interface of the report editor may enable a user to define reporting criteria for the report using the fields associated with the selected data model object, and the events used to generate the report may be constrained to the events that match, or otherwise satisfy, the search constraints of the selected data model object.

(234) The selection of a data model object for use in driving a report generation may be facilitated

by a data model object selection interface. FIG. 9 illustrates an example interactive data model selection graphical user interface **900** of a report editor that displays a listing of available data models **901**. The user may select one of the data models **902**.

(235) FIG. 10 illustrates an example data model object selection graphical user interface **1000** that displays available data objects **1001** for the selected data object model **902**. The user may select one of the displayed data model objects **1002** for use in driving the report generation process.

(236) Once a data model object is selected by the user, a user interface screen **1100** shown in FIG. 11A may display an interactive listing of automatic field identification options **1101** based on the selected data model object. For example, a user may select one of the three illustrated options (e.g., the “All Fields” option **1102**, the “Selected Fields” option **1103**, or the “Coverage” option (e.g., fields with at least a specified % of coverage) **1104**). If the user selects the “All Fields” option **1102**, all of the fields identified from the events that were returned in response to an initial search query may be selected. That is, for example, all of the fields of the identified data model object fields may be selected. If the user selects the “Selected Fields” option **1103**, only the fields from the fields of the identified data model object fields that are selected by the user may be used. If the user selects the “Coverage” option **1104**, only the fields of the identified data model object fields meeting a specified coverage criteria may be selected. A percent coverage may refer to the percentage of events returned by the initial search query that a given field appears in. Thus, for example, if an object dataset includes 10,000 events returned in response to an initial search query, and the “avg\_age” field appears in 854 of those 10,000 events, then the “avg\_age” field would have a coverage of 8.54% for that object dataset. If, for example, the user selects the “Coverage” option and specifies a coverage value of 2%, only fields having a coverage value equal to or greater than 2% may be selected. The number of fields corresponding to each selectable option may be displayed in association with each option. For example, “97” displayed next to the “All Fields” option **1102** indicates that 97 fields will be selected if the “All Fields” option is selected. The “3” displayed next to the “Selected Fields” option **1103** indicates that 3 of the 97 fields will be selected if the “Selected Fields” option is selected. The “49” displayed next to the “Coverage” option **1104** indicates that 49 of the 97 fields (e.g., the 49 fields having a coverage of 2% or greater) will be selected if the “Coverage” option is selected. The number of fields corresponding to the “Coverage” option may be dynamically updated based on the specified percent of coverage.

(237) FIG. 11B illustrates an example graphical user interface screen **1105** displaying the reporting application's “Report Editor” page. The screen may display interactive elements for defining various elements of a report. For example, the page includes a “Filters” element **1106**, a “Split Rows” element **1107**, a “Split Columns” element **1108**, and a “Column Values” element **1109**. The page may include a list of search results **1111**. In this example, the Split Rows element **1107** is expanded, revealing a listing of fields **1110** that can be used to define additional criteria (e.g., reporting criteria). The listing of fields **1110** may correspond to the selected fields. That is, the listing of fields **1110** may list only the fields previously selected, either automatically and/or manually by a user. FIG. 11C illustrates a formatting dialogue **1112** that may be displayed upon selecting a field from the listing of fields **1110**. The dialogue can be used to format the display of the results of the selection (e.g., label the column for the selected field to be displayed as “component”).

(238) FIG. 11D illustrates an example graphical user interface screen **1105** including a table of results **1113** based on the selected criteria including splitting the rows by the “component” field. A column **1114** having an associated count for each component listed in the table may be displayed that indicates an aggregate count of the number of times that the particular field-value pair (e.g., the value in a row for a particular field, such as the value “BucketMover” for the field “component”) occurs in the set of events responsive to the initial search query.

(239) FIG. 12 illustrates an example graphical user interface screen **1200** that allows the user to filter search results and to perform statistical analysis on values extracted from specific fields in the

set of events. In this example, the top ten product names ranked by price are selected as a filter **1201** that causes the display of the ten most popular products sorted by price. Each row is displayed by product name and price **1202**. This results in each product displayed in a column labeled “product name” along with an associated price in a column labeled “price” **1206**. Statistical analysis of other fields in the events associated with the ten most popular products have been specified as column values **1203**. A count of the number of successful purchases for each product is displayed in column **1204**. These statistics may be produced by filtering the search results by the product name, finding all occurrences of a successful purchase in a field within the events and generating a total of the number of occurrences. A sum of the total sales is displayed in column **1205**, which is a result of the multiplication of the price and the number of successful purchases for each product.

(240) The reporting application allows the user to create graphical visualizations of the statistics generated for a report. For example, FIG. **13** illustrates an example graphical user interface **1300** that displays a set of components and associated statistics **1301**. The reporting application allows the user to select a visualization of the statistics in a graph (e.g., bar chart, scatter plot, area chart, line chart, pie chart, radial gauge, marker gauge, filler gauge, etc.), where the format of the graph may be selected using the user interface controls **1302** along the left panel of the user interface **1300**. FIG. **14** illustrates an example of a bar chart visualization **1400** of an aspect of the statistical data **1301**. FIG. **15** illustrates a scatter plot visualization **1500** of an aspect of the statistical data **1301**.

(241) The above-described system provides significant flexibility by enabling a user to analyze massive quantities of minimally-processed data “on the fly” at search time using a late-binding schema, instead of storing pre-specified portions of the data in a database at ingestion time. This flexibility enables a user to see valuable insights, correlate data, and perform subsequent queries to examine interesting aspects of the data that may not have been apparent at ingestion time.

(242) However, performing extraction and analysis operations at search time can involve a large amount of data and require a large number of computational operations, which can cause delays in processing the queries. Advantageously, the data intake and query system also employs a number of unique acceleration techniques that have been developed to speed up analysis operations performed at search time. These techniques include: (1) performing search operations in parallel across multiple indexers; (2) using a keyword index; (3) using a high performance analytics store; and (4) accelerating the process of generating reports. These novel techniques are described in more detail below.

(243) To facilitate faster query processing, a query can be structured such that multiple indexers perform the query in parallel, while aggregation of search results from the multiple indexers is performed locally at the search head. For example, FIG. **16** is an example search query received from a client and executed by search peers, in accordance with example embodiments. FIG. **16** illustrates how a search query **1602** received from a client at a search head **210** can split into two phases, including: (1) subtasks **1604** (e.g., data retrieval or simple filtering) that may be performed in parallel by indexers **206** for execution, and (2) a search results aggregation operation **1606** to be executed by the search head when the results are ultimately collected from the indexers.

(244) During operation, upon receiving search query **1602**, a search head **210** determines that a portion of the operations involved with the search query may be performed locally by the search head. The search head modifies search query **1602** by substituting “stats” (create aggregate statistics over results sets received from the indexers at the search head) with “prestats” (create statistics by the indexer from local results set) to produce search query **1604**, and then distributes search query **1604** to distributed indexers, which are also referred to as “search peers” or “peer indexers.” Note that search queries may generally specify search criteria or operations to be performed on events that meet the search criteria. Search queries may also specify field names, as well as search criteria for the values in the fields or operations to be performed on the values in the

fields. Moreover, the search head may distribute the full search query to the search peers as illustrated in FIG. 6A, or may alternatively distribute a modified version (e.g., a more restricted version) of the search query to the search peers. In this example, the indexers are responsible for producing the results and sending them to the search head. After the indexers return the results to the search head, the search head aggregates the received results **1606** to form a single search result set. By executing the query in this manner, the system effectively distributes the computational operations across the indexers while minimizing data transfers.

(245) As described above with reference to the flow charts in FIG. 5A and FIG. 6A, data intake and query system **108** can construct and maintain one or more keyword indices to quickly identify events containing specific keywords. This technique can greatly speed up the processing of queries involving specific keywords. As mentioned above, to build a keyword index, an indexer first identifies a set of keywords. Then, the indexer includes the identified keywords in an index, which associates each stored keyword with references to events containing that keyword, or to locations within events where that keyword is located. When an indexer subsequently receives a keyword-based query, the indexer can access the keyword index to quickly identify events containing the keyword.

(246) To speed up certain types of queries, some embodiments of system **108** create a high performance analytics store, which is referred to as a “summarization table,” that contains entries for specific field-value pairs. Each of these entries keeps track of instances of a specific value in a specific field in the events and includes references to events containing the specific value in the specific field. For example, an example entry in a summarization table can keep track of occurrences of the value “94107” in a “ZIP code” field of a set of events and the entry includes references to all of the events that contain the value “94107” in the ZIP code field. This optimization technique enables the system to quickly process queries that seek to determine how many events have a particular value for a particular field. To this end, the system can examine the entry in the summarization table to count instances of the specific value in the field without having to go through the individual events or perform data extractions at search time. Also, if the system needs to process all events that have a specific field-value combination, the system can use the references in the summarization table entry to directly access the events to extract further information without having to search all of the events to find the specific field-value combination at search time.

(247) In some embodiments, the system maintains a separate summarization table for each of the above-described time-specific buckets that stores events for a specific time range. A bucket-specific summarization table includes entries for specific field-value combinations that occur in events in the specific bucket. Alternatively, the system can maintain a separate summarization table for each indexer. The indexer-specific summarization table includes entries for the events in a data store that are managed by the specific indexer. Indexer-specific summarization tables may also be bucket-specific.

(248) The summarization table can be populated by running a periodic query that scans a set of events to find instances of a specific field-value combination, or alternatively instances of all field-value combinations for a specific field. A periodic query can be initiated by a user, or can be scheduled to occur automatically at specific time intervals. A periodic query can also be automatically launched in response to a query that asks for a specific field-value combination.

(249) In some cases, when the summarization tables may not cover all of the events that are relevant to a query, the system can use the summarization tables to obtain partial results for the events that are covered by summarization tables, but may also have to search through other events that are not covered by the summarization tables to produce additional results. These additional results can then be combined with the partial results to produce a final set of results for the query. The summarization table and associated techniques are described in more detail in U.S. Pat. No. 8,682,925, entitled “DISTRIBUTED HIGH PERFORMANCE ANALYTICS STORE”, issued on



25 Mar. 2014, U.S. Pat No. 9,128,985, entitled “SUPPLEMENTING A HIGH PERFORMANCE ANALYTICS STORE WITH EVALUATION OF INDIVIDUAL EVENTS TO RESPOND TO AN EVENT QUERY”, issued on 8 Sep. 2015, and U.S. patent application Ser. No. 14/815,973, entitled “GENERATING AND STORING SUMMARIZATION TABLES FOR SETS OF SEARCHABLE EVENTS”, filed on 1 Aug. 2015, each of which is hereby incorporated by reference in its entirety for all purposes.

(250) To speed up certain types of queries, e.g., frequently encountered queries or computationally intensive queries, some embodiments of system **108** create a high performance analytics store, which is referred to as a “summarization table,” (also referred to as a “lexicon” or “inverted index”) that contains entries for specific field-value pairs. Each of these entries keeps track of instances of a specific value in a specific field in the event data and includes references to events containing the specific value in the specific field. For example, an example entry in an inverted index can keep track of occurrences of the value “94107” in a “ZIP code” field of a set of events and the entry includes references to all of the events that contain the value “94107” in the ZIP code field.

Creating the inverted index data structure avoids needing to incur the computational overhead each time a statistical query needs to be run on a frequently encountered field-value pair. In order to expedite queries, in most embodiments, the search engine will employ the inverted index separate from the raw record data store to generate responses to the received queries.

(251) Note that the term “summarization table” or “inverted index” as used herein is a data structure that may be generated by an indexer that includes at least field names and field values that have been extracted and/or indexed from event records. An inverted index may also include reference values that point to the location(s) in the field searchable data store where the event records that include the field may be found. Also, an inverted index may be stored using well-known compression techniques to reduce its storage size.

(252) Further, note that the term “reference value” (also referred to as a “posting value”) as used herein is a value that references the location of a source record in the field searchable data store. In some embodiments, the reference value may include additional information about each record, such as timestamps, record size, meta-data, or the like. Each reference value may be a unique identifier which may be used to access the event data directly in the field searchable data store. In some embodiments, the reference values may be ordered based on each event record's timestamp. For example, if numbers are used as identifiers, they may be sorted so event records having a later timestamp always have a lower valued identifier than event records with an earlier timestamp, or vice-versa. Reference values are often included in inverted indexes for retrieving and/or identifying event records.

(253) In one or more embodiments, an inverted index is generated in response to a user-initiated collection query. The term “collection query” as used herein refers to queries that include commands that generate summarization information and inverted indexes (or summarization tables) from event records stored in the field searchable data store.

(254) Note that a collection query is a special type of query that can be user-generated and is used to create an inverted index. A collection query is not the same as a query that is used to call up or invoke a pre-existing inverted index. In one or more embodiments, a query can comprise an initial step that calls up a pre-generated inverted index on which further filtering and processing can be performed. For example, referring back to FIG. 6B, a set of events can be generated at block **640** by either using a “collection” query to create a new inverted index or by calling up a pre-generated inverted index. A query with several pipelined steps will start with a pre-generated index to accelerate the query.

(255) FIG. 7C illustrates the manner in which an inverted index is created and used in accordance with the disclosed embodiments. As shown in FIG. 7C, an inverted index **722** can be created in response to a user-initiated collection query using the event data **723** stored in the raw record data store. For example, a non-limiting example of a collection query may include “collect

clientip=127.0.0.1" which may result in an inverted index **722** being generated from the event data **723** as shown in FIG. 7C. Each entry in inverted index **722** includes an event reference value that references the location of a source record in the field searchable data store. The reference value may be used to access the original event record directly from the field searchable data store.

(256) In one or more embodiments, if one or more of the queries is a collection query, the responsive indexers may generate summarization information based on the fields of the event records located in the field searchable data store. In at least one of the various embodiments, one or more of the fields used in the summarization information may be listed in the collection query and/or they may be determined based on terms included in the collection query. For example, a collection query may include an explicit list of fields to summarize or, in at least one of the various embodiments, a collection query may include terms or expressions that explicitly define the fields, e.g., using regex rules. In FIG. 7C, prior to running the collection query that generates the inverted index **722**, the field name "clientip" may need to be defined in a configuration file by specifying the "access\_combined" source type and a regular expression rule to parse out the client IP address. Alternatively, the collection query may contain an explicit definition for the field name "clientip" which may obviate the need to reference the configuration file at search time.

(257) In one or more embodiments, collection queries may be saved and scheduled to run periodically. These scheduled collection queries may periodically update the summarization information corresponding to the query. For example, if the collection query that generates inverted index **722** is scheduled to run periodically, one or more indexers would periodically search through the relevant buckets to update inverted index **722** with event data for any new events with the "clientip" value of "127.0.0.1."

(258) In some embodiments, the inverted indexes that include fields, values, and reference value (e.g., inverted index **722**) for event records may be included in the summarization information provided to the user. In other embodiments, a user may not be interested in specific fields and values contained in the inverted index, but may need to perform a statistical query on the data in the inverted index. For example, referencing the example of FIG. 7C rather than viewing the fields within summarization table **722**, a user may want to generate a count of all client requests from IP address "127.0.0.1." In this case, the search engine would simply return a result of "4" rather than including details about the inverted index **722** in the information provided to the user.

(259) The pipelined search language, e.g., SPL of the SPLUNK® ENTERPRISE system can be used to pipe the contents of an inverted index to a statistical query using the "stats" command for example. A "stats" query refers to queries that generate result sets that may produce aggregate and statistical results from event records, e.g., average, mean, max, min, rms, etc. Where sufficient information is available in an inverted index, a "stats" query may generate their result sets rapidly from the summarization information available in the inverted index rather than directly scanning event records. For example, the contents of inverted index **722** can be pipelined to a stats query, e.g., a "count" function that counts the number of entries in the inverted index and returns a value of "4." In this way, inverted indexes may enable various stats queries to be performed absent scanning or search the event records. Accordingly, this optimization technique enables the system to quickly process queries that seek to determine how many events have a particular value for a particular field. To this end, the system can examine the entry in the inverted index to count instances of the specific value in the field without having to go through the individual events or perform data extractions at search time.

(260) In some embodiments, the system maintains a separate inverted index for each of the above-described time-specific buckets that stores events for a specific time range. A bucket-specific inverted index includes entries for specific field-value combinations that occur in events in the specific bucket. Alternatively, the system can maintain a separate inverted index for each indexer. The indexer-specific inverted index includes entries for the events in a data store that are managed by the specific indexer. Indexer-specific inverted indexes may also be bucket-specific. In at least

one or more embodiments, if one or more of the queries is a stats query, each indexer may generate a partial result set from previously generated summarization information. The partial result sets may be returned to the search head that received the query and combined into a single result set for the query

(261) As mentioned above, the inverted index can be populated by running a periodic query that scans a set of events to find instances of a specific field-value combination, or alternatively instances of all field-value combinations for a specific field. A periodic query can be initiated by a user, or can be scheduled to occur automatically at specific time intervals. A periodic query can also be automatically launched in response to a query that asks for a specific field-value combination. In some embodiments, if summarization information is absent from an indexer that includes responsive event records, further actions may be taken, such as, the summarization information may be generated on the fly, warnings may be provided the user, the collection query operation may be halted, the absence of summarization information may be ignored, or the like, or combination thereof.

(262) In one or more embodiments, an inverted index may be set up to update continually. For example, the query may ask for the inverted index to update its result periodically, e.g., every hour. In such instances, the inverted index may be a dynamic data structure that is regularly updated to include information regarding incoming events.

(263) In some cases, e.g., where a query is executed before an inverted index updates, when the inverted index may not cover all of the events that are relevant to a query, the system can use the inverted index to obtain partial results for the events that are covered by inverted index, but may also have to search through other events that are not covered by the inverted index to produce additional results on the fly. In other words, an indexer would need to search through event data on the data store to supplement the partial results. These additional results can then be combined with the partial results to produce a final set of results for the query. Note that in typical instances where an inverted index is not completely up to date, the number of events that an indexer would need to search through to supplement the results from the inverted index would be relatively small. In other words, the search to get the most recent results can be quick and efficient because only a small number of event records will be searched through to supplement the information from the inverted index. The inverted index and associated techniques are described in more detail in U.S. Pat. No. 8,682,925, entitled “DISTRIBUTED HIGH PERFORMANCE ANALYTICS STORE”, issued on 25 Mar. 2014, U.S. Pat. No. 9,128,985, entitled “SUPPLEMENTING A HIGH PERFORMANCE ANALYTICS STORE WITH EVALUATION OF INDIVIDUAL EVENTS TO RESPOND TO AN EVENT QUERY”, filed on 31 Jan. 2014, and U.S. patent application Ser. No. 14/815,973, entitled “STORAGE MEDIUM AND CONTROL DEVICE”, filed on 21 Feb. 2014, each of which is hereby incorporated by reference in its entirety.

(264) In one or more embodiments, if the system needs to process all events that have a specific field-value combination, the system can use the references in the inverted index entry to directly access the events to extract further information without having to search all of the events to find the specific field-value combination at search time. In other words, the system can use the reference values to locate the associated event data in the field searchable data store and extract further information from those events, e.g., extract further field values from the events for purposes of filtering or processing or both.

(265) The information extracted from the event data using the reference values can be directed for further filtering or processing in a query using the pipeline search language. The pipelined search language will, in one embodiment, include syntax that can direct the initial filtering step in a query to an inverted index. In one embodiment, a user would include syntax in the query that explicitly directs the initial searching or filtering step to the inverted index.

(266) Referencing the example in FIG. 7C, if the user determines that she needs the user id fields associated with the client requests from IP address “127.0.0.1,” instead of incurring the

computational overhead of performing a brand new search or re-generating the inverted index with an additional field, the user can generate a query that explicitly directs or pipes the contents of the already generated inverted index 722 to another filtering step requesting the user ids for the entries in inverted index 722 where the server response time is greater than “0.0900” microseconds. The search engine would use the reference values stored in inverted index 722 to retrieve the event data from the field searchable data store, filter the results based on the “response time” field values and, further, extract the user id field from the resulting event data to return to the user. In the present instance, the user ids “frank” and “carlos” would be returned to the user from the generated results table 722.

(267) In one embodiment, the same methodology can be used to pipe the contents of the inverted index to a processing step. In other words, the user is able to use the inverted index to efficiently and quickly perform aggregate functions on field values that were not part of the initially generated inverted index. For example, a user may want to determine an average object size (size of the requested gif) requested by clients from IP address “127.0.0.1.” In this case, the search engine would again use the reference values stored in inverted index 722 to retrieve the event data from the field searchable data store and, further, extract the object size field values from the associated events 731, 732, 733 and 734. Once, the corresponding object sizes have been extracted (i.e. 2326, 2900, 2920, and 5000), the average can be computed and returned to the user.

(268) In one embodiment, instead of explicitly invoking the inverted index in a user-generated query, e.g., by the use of special commands or syntax, the SPLUNK® ENTERPRISE system can be configured to automatically determine if any prior-generated inverted index can be used to expedite a user query. For example, the user's query may request the average object size (size of the requested gif) requested by clients from IP address “127.0.0.1.” without any reference to or use of inverted index 722. The search engine, in this case, would automatically determine that an inverted index 722 already exists in the system that could expedite this query. In one embodiment, prior to running any search comprising a field-value pair, for example, a search engine may search through all the existing inverted indexes to determine if a pre-generated inverted index could be used to expedite the search comprising the field-value pair. Accordingly, the search engine would automatically use the pre-generated inverted index, e.g., index 722 to generate the results without any user-involvement that directs the use of the index.

(269) Using the reference values in an inverted index to be able to directly access the event data in the field searchable data store and extract further information from the associated event data for further filtering and processing is highly advantageous because it avoids incurring the computation overhead of regenerating the inverted index with additional fields or performing a new search.

(270) The data intake and query system includes one or more forwarders that receive raw machine data from a variety of input data sources, and one or more indexers that process and store the data in one or more data stores. By distributing events among the indexers and data stores, the indexers can analyze events for a query in parallel. In one or more embodiments, a multiple indexer implementation of the search system would maintain a separate and respective inverted index for each of the above-described time-specific buckets that stores events for a specific time range. A bucket-specific inverted index includes entries for specific field-value combinations that occur in events in the specific bucket. As explained above, a search head would be able to correlate and synthesize data from across the various buckets and indexers.

(271) This feature advantageously expedites searches because instead of performing a computationally intensive search in a centrally located inverted index that catalogues all the relevant events, an indexer is able to directly search an inverted index stored in a bucket associated with the time-range specified in the query. This allows the search to be performed in parallel across the various indexers. Further, if the query requests further filtering or processing to be conducted on the event data referenced by the locally stored bucket-specific inverted index, the indexer is able to simply access the event records stored in the associated bucket for further filtering and

processing instead of needing to access a central repository of event records, which would dramatically add to the computational overhead.

(272) In one embodiment, there may be multiple buckets associated with the time-range specified in a query. If the query is directed to an inverted index, or if the search engine automatically determines that using an inverted index would expedite the processing of the query, the indexers will search through each of the inverted indexes associated with the buckets for the specified time-range. This feature allows the High Performance Analytics Store to be scaled easily.

(273) In certain instances, where a query is executed before a bucket-specific inverted index updates, when the bucket-specific inverted index may not cover all of the events that are relevant to a query, the system can use the bucket-specific inverted index to obtain partial results for the events that are covered by bucket-specific inverted index, but may also have to search through the event data in the bucket associated with the bucket-specific inverted index to produce additional results on the fly. In other words, an indexer would need to search through event data stored in the bucket (that was not yet processed by the indexer for the corresponding inverted index) to supplement the partial results from the bucket-specific inverted index.

(274) FIG. 7D presents a flowchart illustrating how an inverted index in a pipelined search query can be used to determine a set of event data that can be further limited by filtering or processing in accordance with the disclosed embodiments.

(275) At block **742**, a query is received by a data intake and query system. In some embodiments, the query can be received as a user generated query entered into search bar of a graphical user search interface. The search interface also includes a time range control element that enables specification of a time range for the query.

(276) At block **744**, an inverted index is retrieved. Note, that the inverted index can be retrieved in response to an explicit user search command inputted as part of the user generated query.

Alternatively, the search engine can be configured to automatically use an inverted index if it determines that using the inverted index would expedite the servicing of the user generated query. Each of the entries in an inverted index keeps track of instances of a specific value in a specific field in the event data and includes references to events containing the specific value in the specific field. In order to expedite queries, in most embodiments, the search engine will employ the inverted index separate from the raw record data store to generate responses to the received queries.

(277) At block **746**, the query engine determines if the query contains further filtering and processing steps. If the query contains no further commands, then, in one embodiment, summarization information can be provided to the user at block **754**.

(278) If, however, the query does contain further filtering and processing commands, then at block **750**, the query engine determines if the commands relate to further filtering or processing of the data extracted as part of the inverted index or whether the commands are directed to using the inverted index as an initial filtering step to further filter and process event data referenced by the entries in the inverted index. If the query can be completed using data already in the generated inverted index, then the further filtering or processing steps, e.g., a “count” number of records function, “average” number of records per hour etc. are performed and the results are provided to the user at block **752**.

(279) If, however, the query references fields that are not extracted in the inverted index, then the indexers will access event data pointed to by the reference values in the inverted index to retrieve any further information required at block **756**. Subsequently, any further filtering or processing steps are performed on the fields extracted directly from the event data and the results are provided to the user at step **758**.

(280) In some embodiments, a data server system such as the data intake and query system can accelerate the process of periodically generating updated reports based on query results. To accelerate this process, a summarization engine automatically examines the query to determine whether generation of updated reports can be accelerated by creating intermediate summaries. If

reports can be accelerated, the summarization engine periodically generates a summary covering data obtained during a latest non-overlapping time period. For example, where the query seeks events meeting a specified criteria, a summary for the time period includes only events within the time period that meet the specified criteria. Similarly, if the query seeks statistics calculated from the events, such as the number of events that match the specified criteria, then the summary for the time period includes the number of events in the period that match the specified criteria.

(281) In addition to the creation of the summaries, the summarization engine schedules the periodic updating of the report associated with the query. During each scheduled report update, the query engine determines whether intermediate summaries have been generated covering portions of the time period covered by the report update. If so, then the report is generated based on the information contained in the summaries. Also, if additional event data has been received and has not yet been summarized, and is required to generate the complete report, the query can be run on these additional events. Then, the results returned by this query on the additional events, along with the partial results obtained from the intermediate summaries, can be combined to generate the updated report. This process is repeated each time the report is updated. Alternatively, if the system stores events in buckets covering specific time ranges, then the summaries can be generated on a bucket-by-bucket basis. Note that producing intermediate summaries can save the work involved in re-running the query for previous time periods, so advantageously only the newer events need to be processed while generating an updated report. These report acceleration techniques are described in more detail in U.S. Pat. No. 8,589,403, entitled "COMPRESSED JOURNALING IN EVENT TRACKING FILES FOR METADATA RECOVERY AND REPLICATION," issued on 19 Nov. 2013, U.S. Pat. No. 8,412,696, entitled "REAL TIME SEARCHING AND REPORTING," issued on 2 Apr. 2011, and U.S. Pat. Nos. 8,589,375 and 8,589,432, both also entitled "REAL TIME SEARCHING AND REPORTING," both issued on 19 Nov. 2013, each of which is hereby incorporated by reference in its entirety for all purposes.

(282) The data intake and query system provides various schemas, dashboards, and visualizations that simplify developers' tasks to create applications with additional capabilities. One such application is an enterprise security application, such as SPLUNK® ENTERPRISE SECURITY, which performs monitoring and alerting operations and includes analytics to facilitate identifying both known and unknown security threats based on large volumes of data stored by the data intake and query system. The enterprise security application provides the security practitioner with visibility into security-relevant threats found in the enterprise infrastructure by capturing, monitoring, and reporting on data from enterprise security devices, systems, and applications. Through the use of the data intake and query system searching and reporting capabilities, the enterprise security application provides a top-down and bottom-up view of an organization's security posture.

(283) The enterprise security application leverages the data intake and query system search-time normalization techniques, saved searches, and correlation searches to provide visibility into security-relevant threats and activity and generate notable events for tracking. The enterprise security application enables the security practitioner to investigate and explore the data to find new or unknown threats that do not follow signature-based patterns.

(284) Conventional Security Information and Event Management (SIEM) systems lack the infrastructure to effectively store and analyze large volumes of security-related data. Traditional SIEM systems typically use fixed schemas to extract data from pre-defined security-related fields at data ingestion time and store the extracted data in a relational database. This traditional data extraction process (and associated reduction in data size) that occurs at data ingestion time inevitably hampers future incident investigations that may need original data to determine the root cause of a security issue, or to detect the onset of an impending security threat.

(285) In contrast, the enterprise security application system stores large volumes of minimally-processed security-related data at ingestion time for later retrieval and analysis at search time when

a live security threat is being investigated. To facilitate this data retrieval process, the enterprise security application provides pre-specified schemas for extracting relevant values from the different types of security-related events and enables a user to define such schemas.

(286) The enterprise security application can process many types of security-related information. In general, this security-related information can include any information that can be used to identify security threats. For example, the security-related information can include network-related information, such as IP addresses, domain names, asset identifiers, network traffic volume, uniform resource locator strings, and source addresses. The process of detecting security threats for network-related information is further described in U.S. Pat. No. 8,826,434, entitled “SECURITY THREAT DETECTION BASED ON INDICATIONS IN BIG DATA OF ACCESS TO NEWLY REGISTERED DOMAINS”, issued on 2 Sep. 2014, U.S. Pat. No. 9,215,240, entitled “INVESTIGATIVE AND DYNAMIC DETECTION OF POTENTIAL SECURITY-THREAT INDICATORS FROM EVENTS IN BIG DATA”, issued on 15 Dec. 2015, U.S. Pat. No. 9,173,801, entitled “GRAPHIC DISPLAY OF SECURITY THREATS BASED ON INDICATIONS OF ACCESS TO NEWLY REGISTERED DOMAINS”, issued on 3 Nov. 2015, U.S. Pat. No. 9,248,068, entitled “SECURITY THREAT DETECTION OF NEWLY REGISTERED DOMAINS”, issued on 2 Feb. 2016, U.S. Pat. No. 9,426,172, entitled “SECURITY THREAT DETECTION USING DOMAIN NAME ACCESSES”, issued on 23 Aug. 2016, and U.S. Pat. No. 9,432,396, entitled “SECURITY THREAT DETECTION USING DOMAIN NAME REGISTRATIONS”, issued on 30 Aug. 2016, each of which is hereby incorporated by reference in its entirety for all purposes. Security-related information can also include malware infection data and system configuration information, as well as access control information, such as login/logout information and access failure notifications. The security-related information can originate from various sources within a data center, such as hosts, virtual machines, storage devices and sensors. The security-related information can also originate from various sources in a network, such as routers, switches, email servers, proxy servers, gateways, firewalls and intrusion-detection systems.

(287) During operation, the enterprise security application facilitates detecting “notable events” that are likely to indicate a security threat. A notable event represents one or more anomalous incidents, the occurrence of which can be identified based on one or more events (e.g., time stamped portions of raw machine data) fulfilling pre-specified and/or dynamically-determined (e.g., based on machine-learning) criteria defined for that notable event. Examples of notable events include the repeated occurrence of an abnormal spike in network usage over a period of time, a single occurrence of unauthorized access to system, a host communicating with a server on a known threat list, and the like. These notable events can be detected in a number of ways, such as: (1) a user can notice a correlation in events and can manually identify that a corresponding group of one or more events amounts to a notable event; or (2) a user can define a “correlation search” specifying criteria for a notable event, and every time one or more events satisfy the criteria, the application can indicate that the one or more events correspond to a notable event; and the like. A user can alternatively select a pre-defined correlation search provided by the application. Note that correlation searches can be run continuously or at regular intervals (e.g., every hour) to search for notable events. Upon detection, notable events can be stored in a dedicated “notable events index,” which can be subsequently accessed to generate various visualizations containing security-related information. Also, alerts can be generated to notify system operators when important notable events are discovered.

(288) The enterprise security application provides various visualizations to aid in discovering security threats, such as a “key indicators view” that enables a user to view security metrics, such as counts of different types of notable events. For example, FIG. 17A illustrates an example key indicators view **1700** that comprises a dashboard, which can display a value **1701**, for various security-related metrics, such as malware infections **1702**. It can also display a change in a metric value **1703**, which indicates that the number of malware infections increased by 63 during the

preceding interval. Key indicators view **1700** additionally displays a histogram panel **1704** that displays a histogram of notable events organized by urgency values, and a histogram of notable events organized by time intervals. This key indicators view is described in further detail in pending U.S. patent application Ser. No. 13/956,338, entitled “KEY INDICATORS VIEW”, filed on 31 Jul. 2013, and which is hereby incorporated by reference in its entirety for all purposes. (289) These visualizations can also include an “incident review dashboard” that enables a user to view and act on “notable events.” These notable events can include: (1) a single event of high importance, such as any activity from a known web attacker; or (2) multiple events that collectively warrant review, such as a large number of authentication failures on a host followed by a successful authentication. For example, FIG. **17B** illustrates an example incident review dashboard **1710** that includes a set of incident attribute fields **1711** that, for example, enables a user to specify a time range field **1712** for the displayed events. It also includes a timeline **1713** that graphically illustrates the number of incidents that occurred in time intervals over the selected time range. It additionally displays an events list **1714** that enables a user to view a list of all of the notable events that match the criteria in the incident attributes fields **1711**. To facilitate identifying patterns among the notable events, each notable event can be associated with an urgency value (e.g., low, medium, high, critical), which is indicated in the incident review dashboard. The urgency value for a detected event can be determined based on the severity of the event and the priority of the system component associated with the event.

(290) As mentioned above, the data intake and query platform provides various features that simplify the developer's task to create various applications. One such application is a virtual machine monitoring application, such as SPLUNK® APP FOR VMWARE® that provides operational visibility into granular performance metrics, logs, tasks and events, and topology from hosts, virtual machines and virtual centers. It empowers administrators with an accurate real-time picture of the health of the environment, proactively identifying performance and capacity bottlenecks.

(291) Conventional data-center-monitoring systems lack the infrastructure to effectively store and analyze large volumes of machine-generated data, such as performance information and log data obtained from the data center. In conventional data-center-monitoring systems, machine-generated data is typically pre-processed prior to being stored, for example, by extracting pre-specified data items and storing them in a database to facilitate subsequent retrieval and analysis at search time. However, the rest of the data is not saved and discarded during pre-processing.

(292) In contrast, the virtual machine monitoring application stores large volumes of minimally processed machine data, such as performance information and log data, at ingestion time for later retrieval and analysis at search time when a live performance issue is being investigated. In addition to data obtained from various log files, this performance-related information can include values for performance metrics obtained through an application programming interface (API) provided as part of the vSphere Hypervisor™ system distributed by VMware, Inc. of Palo Alto, California. For example, these performance metrics can include: (1) CPU-related performance metrics; (2) disk-related performance metrics; (3) memory-related performance metrics; (4) network-related performance metrics; (5) energy-usage statistics; (6) data-traffic-related performance metrics; (7) overall system availability performance metrics; (8) cluster-related performance metrics; and (9) virtual machine performance statistics. Such performance metrics are described in U.S. patent application Ser. No. 14/167,316, entitled “CORRELATION FOR USER-SELECTED TIME RANGES OF VALUES FOR PERFORMANCE METRICS OF COMPONENTS IN AN INFORMATION-TECHNOLOGY ENVIRONMENT WITH LOG DATA FROM THAT INFORMATION-TECHNOLOGY ENVIRONMENT”, filed on 29 Jan. 2014, and which is hereby incorporated by reference in its entirety for all purposes.

(293) To facilitate retrieving information of interest from performance data and log files, the virtual machine monitoring application provides pre-specified schemas for extracting relevant values from



different types of performance-related events, and also enables a user to define such schemas.

(294) The virtual machine monitoring application additionally provides various visualizations to facilitate detecting and diagnosing the root cause of performance problems. For example, one such visualization is a “proactive monitoring tree” that enables a user to easily view and understand relationships among various factors that affect the performance of a hierarchically structured computing system. This proactive monitoring tree enables a user to easily navigate the hierarchy by selectively expanding nodes representing various entities (e.g., virtual centers or computing clusters) to view performance information for lower-level nodes associated with lower-level entities (e.g., virtual machines or host systems). Example node-expansion operations are illustrated in FIG. 17C, wherein nodes **1733** and **1734** are selectively expanded. Note that nodes **1731-1739** can be displayed using different patterns or colors to represent different performance states, such as a critical state, a warning state, a normal state or an unknown/offline state. The ease of navigation provided by selective expansion in combination with the associated performance-state information enables a user to quickly diagnose the root cause of a performance problem. The proactive monitoring tree is described in further detail in U.S. Pat. No. 9,185,007, entitled “PROACTIVE MONITORING TREE WITH SEVERITY STATE SORTING,” issued on 10 Nov. 2015, and U.S. Pat. No. 9,426,045, also entitled “PROACTIVE MONITORING TREE WITH SEVERITY STATE SORTING,” issued on 23 Aug. 2016, each of which is hereby incorporated by reference in its entirety for all purposes.

(295) The virtual machine monitoring application also provides a user interface that enables a user to select a specific time range and then view heterogeneous data comprising events, log data, and associated performance metrics for the selected time range. For example, the screen illustrated in FIG. 17D displays a listing of recent “tasks and events” and a listing of recent “log entries” for a selected time range above a performance-metric graph for “average CPU core utilization” for the selected time range. Note that a user is able to operate pull-down menus **1742** to selectively display different performance metric graphs for the selected time range. This enables the user to correlate trends in the performance-metric graph with corresponding event and log data to quickly determine the root cause of a performance problem. This user interface is described in more detail in U.S. patent application Ser. No. 14/167,316, entitled “CORRELATION FOR USER-SELECTED TIME RANGES OF VALUES FOR PERFORMANCE METRICS OF COMPONENTS IN AN INFORMATION-TECHNOLOGY ENVIRONMENT WITH LOG DATA FROM THAT INFORMATION-TECHNOLOGY ENVIRONMENT”, filed on 29 Jan. 2014, and which is hereby incorporated by reference in its entirety for all purposes.

(296) As previously mentioned, the data intake and query platform provides various schemas, dashboards and visualizations that make it easy for developers to create applications to provide additional capabilities. One such application is an IT monitoring application, such as SPLUNK® IT SERVICE INTELLIGENCE™, which performs monitoring and alerting operations. The IT monitoring application also includes analytics to help an analyst diagnose the root cause of performance problems based on large volumes of data stored by the data intake and query system as correlated to the various services an IT organization provides (a service-centric view). This differs significantly from conventional IT monitoring systems that lack the infrastructure to effectively store and analyze large volumes of service-related events. Traditional service monitoring systems typically use fixed schemas to extract data from pre-defined fields at data ingestion time, wherein the extracted data is typically stored in a relational database. This data extraction process and associated reduction in data content that occurs at data ingestion time inevitably hampers future investigations, when all of the original data may be needed to determine the root cause of or contributing factors to a service issue.

(297) In contrast, an IT monitoring application system stores large volumes of minimally-processed service-related data at ingestion time for later retrieval and analysis at search time, to perform regular monitoring, or to investigate a service issue. To facilitate this data retrieval process, the IT

monitoring application enables a user to define an IT operations infrastructure from the perspective of the services it provides. In this service-centric approach, a service such as corporate e-mail may be defined in terms of the entities employed to provide the service, such as host machines and network devices. Each entity is defined to include information for identifying all of the events that pertains to the entity, whether produced by the entity itself or by another machine, and considering the many various ways the entity may be identified in machine data (such as by a URL, an IP address, or machine name). The service and entity definitions can organize events around a service so that all of the events pertaining to that service can be easily identified. This capability provides a foundation for the implementation of Key Performance Indicators.

(298) One or more Key Performance Indicators (KPI's) are defined for a service within the IT monitoring application. Each KPI measures an aspect of service performance at a point in time or over a period of time (aspect KPI's). Each KPI is defined by a search query that derives a KPI value from the machine data of events associated with the entities that provide the service. Information in the entity definitions may be used to identify the appropriate events at the time a KPI is defined or whenever a KPI value is being determined. The KPI values derived over time may be stored to build a valuable repository of current and historical performance information for the service, and the repository, itself, may be subject to search query processing. Aggregate KPIs may be defined to provide a measure of service performance calculated from a set of service aspect KPI values; this aggregate may even be taken across defined timeframes and/or across multiple services. A particular service may have an aggregate KPI derived from substantially all of the aspect KPI's of the service to indicate an overall health score for the service.

(299) The IT monitoring application facilitates the production of meaningful aggregate KPI's through a system of KPI thresholds and state values. Different KPI definitions may produce values in different ranges, and so the same value may mean something very different from one KPI definition to another. To address this, the IT monitoring application implements a translation of individual KPI values to a common domain of "state" values. For example, a KPI range of values may be 1-100, or 50-275, while values in the state domain may be 'critical,' 'warning,' 'normal,' and 'informational'. Thresholds associated with a particular KPI definition determine ranges of values for that KPI that correspond to the various state values. In one case, KPI values 95-100 may be set to correspond to 'critical' in the state domain. KPI values from disparate KPI's can be processed uniformly once they are translated into the common state values using the thresholds. For example, "normal 80% of the time" can be applied across various KPI's. To provide meaningful aggregate KPI's, a weighting value can be assigned to each KPI so that its influence on the calculated aggregate KPI value is increased or decreased relative to the other KPI's.

(300) One service in an IT environment often impacts, or is impacted by, another service. The IT monitoring application can reflect these dependencies. For example, a dependency relationship between a corporate e-mail service and a centralized authentication service can be reflected by recording an association between their respective service definitions. The recorded associations establish a service dependency topology that informs the data or selection options presented in a GUI, for example. The service dependency topology is like a "map" showing how services are connected based on their dependencies. The service topology may itself be depicted in a GUI and may be interactive to allow navigation among related services.

(301) Entity definitions in the IT monitoring application can include informational fields that can serve as metadata, implied data fields, or attributed data fields for the events identified by other aspects of the entity definition. Entity definitions in the IT monitoring application can also be created and updated by an import of tabular data (as represented in a CSV, another delimited file, or a search query result set). The import may be GUI-mediated or processed using import parameters from a GUI-based import definition process. Entity definitions in the IT monitoring application can also be associated with a service by means of a service definition rule. Processing the rule results in the matching entity definitions being associated with the service definition. The rule can be

processed at creation time, and thereafter on a scheduled or on-demand basis. This allows dynamic, rule-based updates to the service definition.

(302) During operation, the IT monitoring application can recognize notable events that may indicate a service performance problem or other situation of interest. These notable events can be recognized by a “correlation search” specifying trigger criteria for a notable event: every time KPI values satisfy the criteria, the application indicates a notable event. A severity level for the notable event may also be specified. Furthermore, when trigger criteria are satisfied, the correlation search may additionally or alternatively cause a service ticket to be created in an IT service management (ITSM) system, such as a systems available from ServiceNow, Inc., of Santa Clara, California.

(303) SPLUNK® IT SERVICE INTELLIGENCE™ provides various visualizations built on its service-centric organization of events and the KPI values generated and collected. Visualizations can be particularly useful for monitoring or investigating service performance. The IT monitoring application provides a service monitoring interface suitable as the home page for ongoing IT service monitoring. The interface is appropriate for settings such as desktop use or for a wall-mounted display in a network operations center (NOC). The interface may prominently display a services health section with tiles for the aggregate KPI's indicating overall health for defined services and a general KPI section with tiles for KPI's related to individual service aspects. These tiles may display KPI information in a variety of ways, such as by being colored and ordered according to factors like the KPI state value. They also can be interactive and navigate to visualizations of more detailed KPI information.

(304) The IT monitoring application provides a service-monitoring dashboard visualization based on a user-defined template. The template can include user-selectable widgets of varying types and styles to display KPI information. The content and the appearance of widgets can respond dynamically to changing KPI information. The KPI widgets can appear in conjunction with a background image, user drawing objects, or other visual elements, that depict the IT operations environment, for example. The KPI widgets or other GUI elements can be interactive so as to provide navigation to visualizations of more detailed KPI information.

(305) The IT monitoring application provides a visualization showing detailed time-series information for multiple KPI's in parallel graph lanes. The length of each lane can correspond to a uniform time range, while the width of each lane may be automatically adjusted to fit the displayed KPI data. Data within each lane may be displayed in a user selectable style, such as a line, area, or bar chart. During operation a user may select a position in the time range of the graph lanes to activate lane inspection at that point in time. Lane inspection may display an indicator for the selected time across the graph lanes and display the KPI value associated with that point in time for each of the graph lanes. The visualization may also provide navigation to an interface for defining a correlation search, using information from the visualization to pre-populate the definition.

(306) The IT monitoring application provides a visualization for incident review showing detailed information for notable events. The incident review visualization may also show summary information for the notable events over a time frame, such as an indication of the number of notable events at each of a number of severity levels. The severity level display may be presented as a rainbow chart with the warmest color associated with the highest severity classification. The incident review visualization may also show summary information for the notable events over a time frame, such as the number of notable events occurring within segments of the time frame. The incident review visualization may display a list of notable events within the time frame ordered by any number of factors, such as time or severity. The selection of a particular notable event from the list may display detailed information about that notable event, including an identification of the correlation search that generated the notable event.

(307) The IT monitoring application provides pre-specified schemas for extracting relevant values from the different types of service-related events. It also enables a user to define such schemas.

(308) FIG. 18 is a block diagram of an embodiment of the data processing environment 200

described previously with reference to FIG. 2 that includes a distributed ledger system **1802** as a data source **202** of the data intake and query system **108**, a distributed ledger system monitor **1804** (also referred to herein as monitor **1804**), and a client device **204** to interact with data associated with the data intake and query system **108**. Non-limiting examples of a distributed ledger system **1802** include, but are not limited to, Ethereum, Hyperledger Fabric, Quorum, Guardtime, KSI, etc. (309) The distributed ledger system monitor **1804** can be used to monitor or obtain data associated with the distributed ledger system **1802**. The monitor **1804** can be implemented using one or more computing devices, virtual machines, containers, pods, another virtualization technology, or the like, in communication with one or more nodes **1806** of the distributed ledger system **1802**. For example, in some embodiments, the monitor **1804** can be implemented on the same or across different computing devices as distinct container instances, with each container having access to a subset of the resources of a host computing device (e.g., a subset of the memory or processing time of the processors of the host computing device), but sharing a similar operating system. For example, the monitor **1804** can be implemented as one or more Docker containers, which are managed by an orchestration platform of an isolated execution environment system, such as Kubernetes.

(310) Although illustrated as being distinct from the data intake and query system **108** and distributed ledger system **1802**, it will be understood that in some embodiments, the monitor **1804** can be implemented as part of the data intake and query system **108** and/or distributed ledger system **1802**. For example, the monitor **1804** can be implemented using or on one or more nodes **1806** of the distributed ledger system **1802** and/or be implemented using one or more components of the data intake and query system **108**. In certain embodiments, such as when the distributed ledger system **1802** is implemented using an isolated execution environment system, such as, but not limited to Kubernetes, Docker, etc., the monitor **1804** can be implemented as an isolated execution environment of the isolated execution environment system and/or using an isolated execution environment system that is separate from the isolated execution environment system used to implement the distributed ledger system **1802**.

(311) In some embodiments, the monitor **1804** interfaces with the distributed ledger system **1802** to collect data from one or more components of the distributed ledger system **1802**, such as the nodes **1806**. In certain embodiments, the monitor **1804** can collect different types of data from the distributed ledger system **1802**. In some embodiments, the monitor **1804** receives, from one or more nodes **1806**, distributed ledger transactions, blocks, metrics data, and/or log data.

(312) Although only one monitor **1804** is shown in FIG. 18, it will be understood that multiple monitors can be used to collect data from the distributed ledger system **1802**. In some embodiments, one or more monitors can collect data from each node **1806** (e.g., from each peer node **1806** and/or ordering node **1806**) or a subset of the nodes **1806** (e.g., one or more peer nodes **1806**).

(313) In some embodiments, the log data can be generated in response to one or more activities on a node, such as an error, receipt of a request from another node **1806** or client computing device, or in response to the node **1806** processing a transaction of the distributed ledger system **1802**. The log data can include information about the activity, such as an identification of the error, a transaction identifier corresponding to the transaction being processed and the nature of the processing task, etc. In some embodiments, the log data can correspond to or identify different transactions that are being processed by the nodes **1806**. For example, the log data generated by a peer node **1806** (as will be described herein) can indicate the processing task being applied to a particular proposed transaction (e.g., receive transaction, endorse transaction, validate/invalidate transaction, commit block with transaction to blockchain, read/write the proposed changes of the transaction to the ledger state **1904**, etc.). Similarly, an ordering node **1806** (as will be described herein) can generate log data indicative of activities it is executing relative to a transaction (e.g., receive endorsed transaction, order transaction, add transaction to a block, communicate transaction

to peer nodes **1806** as part of the block, committing transaction to blockchain as part of a block, etc.).

(314) Depending on the implementation of the nodes **1806**, the log data can be stored in a data store of the nodes, and/or converted and stored as part of log data of an isolated execution environment system, etc. For example, if the nodes **1806** are implemented using one or more isolated execution environments, the log data may undergo processing by the isolated execution environment system and stored as part of a log file of the isolated execution environment system. For example, the log data may be wrapped in a JSON wrapper and stored as part of a Docker or Kubernetes log file, etc.

(315) As described herein, the generated metrics can include information about the performance metrics of the node **1806** and/or the distributed ledger system **1802**, such as, but not limited to, (1) CPU-related performance metrics; (2) disk-related performance metrics; (3) memory-related performance metrics; (4) network-related performance metrics; (5) energy-usage statistics; (6) data-traffic-related performance metrics; (7) overall system availability performance metrics; (8) cluster-related performance metrics; and (9) virtual machine performance statistics, etc. In some cases, the metrics are stored in a data store associated with a node **1806**. In some cases, the metrics can include a timestamp corresponding to when the metric was measured/obtained.

(316) The transaction notifications can include information about a block (including its transactions) that is to be committed to a blockchain. In some cases, the transaction notifications can correspond to individual transactions of a block, the entire block, or parts of a transactions, such as the bytecode used as part of a transaction, etc. In some cases, the transaction notifications can include the entire content of a block (e.g., the header portion, body portion, transactions, metadata, etc.), or a summary of information, such as an indication of which transactions of a block were validated/invalidated and/or committed to a blockchain. In certain embodiments, the transaction notifications can be stored in a data store, a publication-subscription (pub-sub) messaging system, or buffer.

(317) The transaction notifications can differ from the log data. For example, the log data can be generated asynchronously as various activities occur on different nodes **1806** (e.g., errors, specific processing tasks, etc.), whereas the transaction notifications can be generated as a result of a block being committed to a blockchain. For example, in some cases, peer nodes **1806** and/or ordering nodes **1806** can generate log data but only peer nodes **1806** can generate transaction notifications. Further, the transaction notifications can differ from log data in that the log data can include unstructured raw machine data, whereas the transaction notifications can include structured data that identifies the block (or portions thereof) that is to be committed to a blockchain or a summary related to transactions of the block that is to be committed (e.g., identification of validated/invalidated transactions). In addition, the transaction notifications can include information about multiple transactions and/or multiple transaction identifiers, whereas the log data may include information about only one transaction and/or only one transaction identifier.

(318) As mentioned, the monitor **1804** can collect any one or any combination of the data generated by the nodes **1806**. In some embodiments, the monitor **1804** is configured to obtain one type of data, such as the transaction notifications. In some such embodiments, the monitor **1804** can interact with a respective node **1806** to obtain the transaction notifications. As described herein, in some cases, the transaction notifications are posted to a pub-sub. As such, the monitor can subscribe to the pub-sub to obtain the relevant transaction notifications. In some cases, a node **1806** is associated with multiple channels and the transaction notifications for the different channels are found on different topics of a pub-sub or on different pub-sub. In these cases, the monitor **1804** can be configured to subscribe to the different topics and/or pub-sub. In this way, the monitor **1804** can collect the relevant transaction notifications from a node **1806**.

(319) In some cases, the monitor **102** processes the transaction notifications. For example, in some cases, portions of the transaction notification, such as the details of the individual transactions, may

be encrypted or encoded. In these examples, the monitor **1804** can decode byte strings to readable UTF8 strings or hex. Further, the transaction notifications may include information about multiple transactions. In some such embodiments, the monitor **102** may parse information about individual transactions and separately communicate the information about individual transactions to the data intake and query system **108** (as well as the entire transaction notification). In certain cases, each communication can include a transaction identifier that identifies the corresponding transaction. The data intake and query system **108** can store the separate communications as individual events. Accordingly, the monitor **1804** can be used to generate multiple events from one transaction notification. In some embodiments, the data intake and query system **108** can store the individual events generated from the transaction notifications in an index that is separate from an index that store metrics data and/or log data.

(320) Furthermore, the monitor **1804** and/or data intake and query system **108** can extract the transaction identifiers from the communications received from the monitor **1804** using one or more regex rules. In some such embodiments, the data intake and query system **108** can store the transaction identifiers in one or more inverted indexes that associate the transaction identifier with the event that includes it. In some cases, the monitor **1804** can extract additional information from the transaction notifications, such as, but not limited to channel information (e.g., the channel associated with the transaction and/or blockchain), node information (e.g., identification of the nodes that endorsed, ordered, and/or validated the transaction), etc. The data intake and query system **108** can store any one or any combination of the extracted information in one or more inverted indexes.

(321) FIG. **19A** is a block diagram illustrating an example of a distributed ledger system **1802** that provides one or more distributed ledgers **1808A-1808F** (generically referred to as ledger(s) **1808**) or blockchains across one or more nodes **1806A-1806F** (generically referred to as node(s) **1806**). The nodes **1806** can communicate via a network **1902**. The network **1902** can be the same as network **104** or a different public or private network.

(322) Each node **1806** can be implemented using individual computing devices, distributed processing systems, servers, isolated execution environments (e.g., containers, virtual machines, etc.), shared computing resources, and so on. In some embodiments, the nodes **1806** can be implemented on the same or as part of different isolated execution environment systems (e.g., as different containers or pods of the same or different Kubernetes cluster or Docker swarm).

(323) In the illustrated embodiment of FIG. **19**, each node **1806** is shown to include a ledger **1808** (which may include more than one ledger), which can be stored across one or more data stores, etc. In some embodiments, the ledger **1808** of each node **1806** can include one or more blockchains, etc. In some cases, the ledgers **1808** of the different nodes **1806** correspond to each other, include the same or matching data entries, or include the same data.

(324) The distributed nodes **1806** can store, maintain and/or update their respective ledger **1808**. Each node **1806** can be configured for storing a version of the distributed ledger **1808** (or a portion thereof), and the distributed ledger **1808** may be updated from time to time with modifications to the ledger **1808** and/or ledger entries, such as insertion of a ledger entry (also referred to herein as a block) or an update of a ledger entry. The distributed ledger system **1802** may be adapted such that, where issues arise with the distributed ledger **1808** (e.g., hash collisions, insertions at the same time, corrupted ledgers/ledger entries), the issues are resolved based at least on issue resolution logic. For example, such logic may be distributed among each of the nodes **1806** and/or their computing systems and can be used to improve or ensure consistency between copies of the ledgers at the different nodes. In some embodiments, issues may arise that can cause a distributed ledger **1808** to “fork” and/or spawn another instance, for example, where a collision cannot be automatically resolved between the nodes. In such cases, the resolution logic can be used to determine when to “fork” or spawn another instance, etc.

(325) It will be understood that each node **1806** can include fewer or more components. For

example, each node **1806** can include processors, buffers, applications, databases, etc. In some cases, the nodes **1806** can include executable instructions or code that when executed by the node **1806** cause the node **1806** to modify a corresponding ledger **1808** or generate a transaction that is to be stored in a block of a blockchain. In some cases, the executable instructions can be bytecode and can be used to implement or execute a smart contract relative to the ledger **1808**.

(326) As described herein, the nodes **1806** can include at least a decentralized set of computing devices and may even include personal computing devices for individuals, and so on. For example, a ledger **1808** may be stored on a large number of publicly available devices, each acting as a “node” for storing a copy of the ledger **1808** (e.g., being collaboratively maintained by anonymous peers on a network). In some embodiments, the ledger **1808** is only stored and maintained on a set of trusted “nodes,” such as on a private network or on the computing systems of authorized users. In some embodiments, a combination and/or a “mix” of both trusted nodes and public nodes may be utilized, with the same and/or different rules being applied to activities performed at each (e.g., a different validation process may be used for untrusted nodes, or simply untrusted nodes may be unable to perform certain activities). In some embodiments, there may be different levels of nodes with differing characteristics and applied logic.

(327) The ledgers **1808**, ledger entries, and/or information stored on the ledger entries may be used to store information received from one or more computing devices. For example, the information may include banking information, other commercial information, smart contracts, etc. Further, the ledger **1808** and ledger entries may utilize encryption technology to facilitate and/or validate digital signatures or the data received from the computing devices.

(328) In some embodiments, the ledger **1808** is publicly accessible. In some embodiments, the ledger **1808** is only accessible to select, authorized nodes having the appropriate permissions. In some embodiments, portions of the ledger **1808** are public and portions of the ledger **1808** are private. When the ledger **1808** is publicly accessible, the ledger **1808** may be adapted to only store information incidental to a transaction or a document relating to a transaction, and may be adapted such that identifiable information is removed but validation information is maintained (e.g., storing a hash value computed from the underlying information). Further, the information stored on the ledger **1808** may be encrypted (non-limiting example: using a public key of a key pair associated with the data intake and query system **108**), redacted, compressed, transformed (e.g., through a one-way transformation or a reversible transformation), and so on.

(329) Each of the one or more nodes **1806** may have, at various times, versions of the ledger **1808**, and the ledger **1808** may be maintained through the propagation of entries and/or updates that may be copied across ledgers **1808**. Ledger entries may contain elements of information (e.g., header information and/or other data). There may be various rules and/or logic involved in activities relating to the ledger entries (e.g., creating, updating, validating, deleting); for example, a majority, supermajority, or unanimous consent between nodes may be enforced as a condition to an activity relating to an entry. In some embodiments, distributed ledgers **1808** are utilized and the ledger entries are adapted to have various linkages to one another such that the integrity of the ledger entries can be reinforced and/or validated. For example, the linkages may include hashes computed based on prior entries in the ledger **1808**, which may be utilized to determine whether a ledger entry is a fraudulent entry by reviewing the correctness of the hash based on performing the hash on information stored on prior entries.

(330) The ledger **1808** may be maintained through, for example, a “distributed network system,” the distributed network system providing decentralized control and storage of the ledger **1808** at the one or more nodes (which may be considered “nodes” of the system). The number of “nodes” may be fixed or vary with time, and increasing or decreasing the number of “nodes” may impact the performance and/or security of the system.

(331) The ledger **1808** copies stored and maintained at each “node” provide cross-validation with one another in the event of conflicts between ledgers **1808**, and various cryptographic and/or

hashing algorithms may be utilized during the generation, updating, deletion, linking, and so on, of ledger entries such that ledger entries have increased resiliency to unauthorized tampering or modification. For example, a blockchain ledger **1808** may be distributed across nodes **1806** and used to track information received from one or more computing devices. The blockchain ledger **1808** may have entries linked to one another using cryptographic records, and entries in the blockchain may be ordered, time stamped, and/or associated with metadata. These and other methods can be used for protection against “double” transfers and unauthorized modification of ledger entries.

(332) FIG. **19B** is a block diagram illustrating another example of a distributed ledger system **1802** that includes different types of nodes **1806**. Specifically, the illustrated example of FIG. **19B** includes four peer nodes **1806A**, **1806C**, **1806D**, **1806F** (generically referred to as peer node(s) **1806**) and two ordering nodes **1806B**, **1806E** (generically referred to as ordering node(s) **1806**). It will be understood that fewer or more nodes can be included as desired. For example, the distributed ledger system **1802** can include only one ordering node **1806** or two or more ordering nodes **1806**. Similarly, the distributed ledger system **1802** can include fewer or more peer nodes **1806** as desired.

(333) As described herein, the peer nodes **1806** and ordering nodes **1806** can be implemented using one or more computing devices, isolated execution environments, etc. In some embodiments, each peer node **1806** and/or ordering node **1806** can be associated with the same or different organization, entity, or user. For example, one company may be associated with or control peer nodes **1806A**, **1806C** and ordering node **1806B**, a second company may be associated with or control peer node **1806D** and ordering node **1806E**, and a third company may be associated with or control peer node **1806F**. A non-limiting example of a distributed ledger system **1802** that includes peer nodes **1806** and ordering nodes **1806** is the Hyperledger Fabric.

(334) For simplicity in describing FIG. **19B**, the peer nodes **1806** and ordering nodes **1806** are described with reference to a common channel that enables private communications/transactions between the illustrated nodes **1806A-1806F**. However, it will be understood that the peer nodes **1806** and ordering nodes **1806** can be associated with multiple channels that each enable private communications/transactions between nodes associated with the channel and/or be associated with multiple consortiums made up of organizations that control the individual nodes **1806**. Further, it will be understood that each peer node **1806** can include one or more peer node ledgers **1808** and/or ledger states **1904** and perform the functions described herein for each channel with which the peer node **1806** is associated. Similarly, each ordering node **1806** can include an ordering node ledger **1808** and perform the functions described herein for each channel with which the ordering node **1806** is associated. In some cases, each channel can include at least one ordering node **1806** and multiple peer nodes **1806**. In certain embodiments, a channel is associated with multiple peer nodes **1806** and only one ordering node **1806**. In certain cases, multiple ordering nodes **1806** can be associated with the same channel.

(335) In the illustrated embodiment of FIG. **19B**, each of the peer nodes **1806A**, **1806C**, **1806D**, **1806E** includes a respective peer node ledger **1808A**, **1808C**, **1808D**, **1808F** (generically referred to as peer node ledger(s) **1808**) and a respective ledger state **1904A**, **1904C**, **1904D**, **1904E** (generically referred to as ledger state(s) **1904**), and can be used to receive proposed transactions from a client computing device (not shown), endorse transactions, communicate endorsed transactions to a client computing device or ordering node **1806**, validate transactions of a block, commit blocks to a respective peer node ledger **1808**, and/or update a respective ledger state **1904**.

(336) Similar to the description of ledgers **1808** with reference to FIG. **19A**, the peer node ledgers **1808** can include one or more ledgers or blockchains. Further, the peer node ledgers **1808** of the different peer nodes **1806** can correspond to each other, include the same or matching entries, transactions, blocks, blockchains, etc. In some cases, the peer node ledger **1808** can include blocks formed from validated transactions, but may exclude invalidated transactions. In certain



embodiments, the peer node ledgers **1808** can include blocks formed from validated and invalidated (or failed) transactions. In certain embodiments, such as embodiments in which an ordering node **1806** maintains an ordering node ledger **1808**, the peer node ledgers **1808** can correspond to or match the ordering node ledgers **1808** of the ordering nodes **1806** and/or be different. For example, in some cases, the ordering node ledgers **1808** can include all endorsed transactions, regardless of whether they are validated and the peer node ledgers **1808** can include endorsed and validated transactions but not endorsed and invalidated or failed transactions. In certain embodiments, the peer node ledgers **1808** can include one ledger or blockchain that matches the ordering node ledger **1808** and another ledger that does not match the ordering node ledger **1808**.

(337) In some cases, the peer node ledger **1808** is generated based on blocks received from an ordering node **1806**. For example, the peer node **1806** can review the transactions of a received block and, if a transaction is validated, can include the transaction as part of a block for the peer node ledger **1808**. Accordingly, in certain embodiments a block of a peer node **1806** may have fewer transactions (or none) compared to a corresponding block received from the ordering node **1806** and/or found in the ordering node ledger **1808**.

(338) As described herein at least with reference to FIG. 20, when a peer node ledger **1808** is implemented as a blockchain, each block of the blockchain can include a header portion (including metadata) and a body portion. The header portion and/or metadata can include a block number (e.g., which block the block is in the blockchain), one or more content identifiers for the current block, a content identifier for a previous block, one or more timestamps (e.g., when block was created, added to the blockchain, etc.), a digital certificate, a public key (of a public-private key pair), a digital signature of the peer node **1806** that added the block to the blockchain, and/or indicators as to whether a transaction of the block is valid/invalid, etc. In addition, in some cases, the header portion can include hashes or content identifiers for individual transactions of a block, and the body portion of a block in the blockchain can include one or more transactions or transaction data associated with a transaction.

(339) In certain embodiments, each transaction can include header information (e.g., bytecode used to generate the transaction, software version, etc.), digital signature of the client computing device that initiated the transaction, a signature or identifier of the endorsing peer nodes **1806** (peer nodes **1806** that signed and/or endorsed the transaction), channel information (which channel the transaction is associated with), a signature or identifier of the ordering node **1806** that ordered the transaction in the block, a proposed change to the peer node ledger **1808**, an expected input/output of the transaction (e.g., the content of the ledger state **1904** before and after the transaction is executed, etc.), etc.

(340) The ledger state **1904** can include one or more key-value pairs reflecting the value or state of the key (of the key-value pair), and can be implemented as a database in one or more data stores of a peer node **1806**. In some embodiments, the ledger state **1904** reflects a current state or value of the keys based on the transactions in the corresponding peer node ledger **1808** or blockchain. As a non-limiting example, if the peer node ledger **1808** reflects transactions (e.g., debits and credits) associated with a particular bank account or other intangible object, the ledger state **1904** can reflect the current value of money in the bank account based on all previous transactions. As another non-limiting example, the ledger state **1904** can reflect a current ownership of a car or other physical object based on previous (validated) transactions associated with the car found in the peer node ledger **1808**. Accordingly, as a peer node **1806** adds a block with one or more transactions to a peer node ledger **1808** or blockchain, the peer node **1806** can update the ledger state **1904** for keys that were altered based on any one or any combination of the (validated) transactions of the block. Similar to the peer node ledgers **1808**, the ledger states **1904** of the different peer nodes **1806** can correspond to each other, include the same or matching key-value pairs, etc.

(341) Although not illustrated, it will be understood that each peer node **1806** can include fewer or more components. For example, as mentioned, each peer node **1806** can include multiple peer node ledgers **1808**, as well as bytecodes, permissions, etc. This information can be stored on one or more data store associated with the peer node **1806**. The permissions can indicate which channels, organizations, or other components, the peer node **1806** is associated with and/or what information the peer node **1806** is allowed to access or edit, etc.

(342) The bytecodes can include executable instructions that the peer node **1806** is to execute and which can generate or be used to endorse or validate transactions for a block of a blockchain. For example, a bytecode can indicate that a peer node **1806** is to read/write information to a ledger state **1904**. A client computing device (not shown) can cause the peer node **1806** to execute the bytecode by providing the peer node **1806** with one or more inputs. For example, if the bytecode is used to reflect the change in ownership of a car, the client computing device can identify the subject car and the identity of the parties involved in the transaction (e.g., buyer and seller). The peer node **1806** can use the bytecode to verify whether the ledger state **1904** includes the identified car and the parties are valid (e.g., identified owner owns the car and buyer is able to purchase the car), etc. Based on the bytecode, the relevant peer nodes **1806** can endorse or validate a transaction that is to be included as part of a block in a blockchain.

(343) FIG. **20** is a block diagram illustrating an embodiment of a blockchain **2000** or distributed ledger that includes blocks that are linked together. The blockchain **2000** can correspond to a peer node blockchain (non-limiting example: include only validated transactions or an indication of valid/invalid transactions) and/or an ordering node blockchain (non-limiting example: include transactions regardless of validation). In the illustrated embodiment, four blocks **2002A**, **2002B**, **2002C**, **2002D** (generically referred to as block(s) **2002**) of the blockchain **2000** are shown, with each block **2002** including a header portion **2004A**, **2004B**, **2004C**, **2004D** (generically referred to as header portion(s) **2004**) and a body portion **2006A**, **2006B**, **2006C**, **2006D** (generically referred to as body portion **2006**). However, it will be understood that each block **2002** can include fewer or more sections, etc. For example, in some embodiments, each block **2002** can include only a body portion **2006** or only a header portions **2004** (e.g., if a peer node **1806** determines that no transactions of a block received from an ordering node **1806** can be validated, the peer node **1806** can generate a block with no transactions). In addition, for simplicity, some details of the blocks **2002** are not shown. For example, additional information can be included in the header portions and/or body portions, etc.

(344) The distributed ledger system **1802** can generate blocks based on various criteria, such as, but not limited to, the passage of a predetermined time interval, the size or amount of data/transactions received, the determination of a solution to a computational puzzle that is determined by a difficulty parameter, the number of block entries (or transactions), or generated content identifiers received, etc. In some embodiments, the distributed ledger system **1802** can generate a block based on a predetermined period of time. For example, the distributed ledger system **1802** can generate a block for the blockchain **2000** once a second, every 10 seconds, once a minute, every 10 min., every hour, etc. In certain embodiments, the distributed ledger system **1802** can generate a block based on the size or amount of data corresponding to one or more transactions. For example, the distributed ledger system **1802** can generate a block for each group of transactions that forms a megabyte or gigabyte of data. In some embodiments, the distributed ledger system **1802** can generate a block based on a node or computing system determining a solution to a computational puzzle that is based on a difficulty parameter. In certain cases, the difficulty parameter changes over time to ensure that blocks are likely to produce on a regular time interval. In some cases, the distributed ledger system can generate a block based on a number of block entries, transactions, or content identifiers. For example, the distributed ledger system **1802** can generate a block for each transaction or each set of 100, 1000, or 1,000,000 transactions, etc. The distributed ledger system **1802** can use any one or any combination of the aforementioned techniques to generate a block.

(345) In the illustrated embodiment, the header portions **2004** include a content identifier (in this example a hash) associated with the previous block (e.g., a hash of the body portion of the previous block) and a content identifier for the current block (e.g., for the body portion of the current block). For example, the header portion **2004B** includes the hash “49vvszj39fjpa,” which corresponds to the hash of the body portion **2006A** and the hash “69yu8q04prb5,” which corresponds to the hash of the body portion **2006B**.

(346) It will be understood that less, different, or more information can be included in the header portions **2004**, as desired. For example, the header portions **2004** can include a nonce, root hash of a Merkle tree, timestamp, difficulty level, software version number, block number indicating the number of blocks in the blockchain that precede the block, etc. The nonce can correspond to the number used to obtain a hash that is smaller than a target hash. For example, in some embodiments, before a group of transactions can be added as a block to the blockchain **2000**, the distributed ledger system **1802** may require that the hash of the content of the block (e.g., the hash of the body portion **2006**) be lower than a threshold number. To meet that criteria, a node **1806** can add a nonce value and hash the combination of the nonce value and the content of the block. If the resulting hash does not meet the size criteria, the node can repeatedly increment the nonce value and take hash again until the threshold is satisfied. The final nonce value can be included in the block.

(347) As another example, the header portions **2004** can include hashes of the entire previous block (header and/or body portion), one or more timestamps (or time range) reflecting the time when the block was started, completed, and/or added to the blockchain **2000**, and/or a difficulty level. In certain cases, the timestamp can correspond to current day and time and/or the amount of time elapsed from a particular time. The difficulty level can, in certain cases, correspond to the size of the hash. In certain cases, a smaller hash can correspond to a higher difficulty level. The root hash can correspond to the hash created based on hashes of multiple transactions, including any hashes of hashes generated by hashing transactions, and so on.

(348) The header portions **2004** can include a content identifier for each transaction included in the body portion **2006**. For example, the header portion **2004** can include a hash of each transaction in the body portion **2006**. In certain embodiments, the header portion **2004** can include a digital certificate, public key, and/or digital signature associated with the peer node **1806** or ordering node **1806** that created it. In some cases, the header portion **2004** (or other metadata) can include an indicator for each transaction indicating whether the transaction was validated by a peer node **1806**. In some embodiments, where the ordering node ledger **1808** and the peer node ledger **1808** are different, the header portion of a block in a peer node ledger **1808** can include an indication of the block in the ordering node ledger to which it relates. For example, if OrderBlock\_12 in an ordering node blockchain includes Transaction\_A that is later invalidated and excluded from a corresponding PeerBlock\_12 in a peer node blockchain, the header portion of the PeerBlock\_12 can include an identifier that identifies OrderBlock\_12 in the ordering node blockchain as including Transaction\_A.

(349) With continued reference to FIG. 20, the body portions **2006** can include one or more block entries for each transaction of the block. In some embodiments, the block entries can be compressed and/or the content of one or more block entries (or all block entries) can be encoded or encrypted using a public key of a key pair associated with the computing device that provided the information for the block entry. In this way, the distributed ledger system **1802** can limit the accessibility of the block entries.

(350) In some embodiments, the block entries can include transaction data, such as but not limited to, a transaction identifier, node signatures (e.g., endorsing/validating peer nodes **1806**, ordering nodes **1806**, etc.), client computing device signatures, proposed ledger changes, expected input/output of the transaction, bytecode identification, inputs into the bytecode, channel information, timestamp of creation, etc. In some cases, each proposed transaction received by a peer node can be assigned (or come with) a transaction identifier or transaction ID. The transaction

identifier can follow the transaction throughout the validation process and/or be included as part of transaction in a block entry of a block.

(351) The digital signatures can include any one or any combination of a digital signature from the client computing device that initiated the proposed transaction, a digital signature corresponding to the peer nodes **1806** that endorsed the transaction, the digital signature of the ordering node(s) **1806** that ordered the transactions in and/or created the block, and/or the digital signature of the peer node **1806** that validated the transaction as part of the block and/or committed the block to the blockchain. In certain cases, the transaction data of a block entry can include the proposed change to the ledger state **1904**, including the proposed key-value pairs before and after the transaction is executed. In certain cases, the transaction data can include an identification of the bytecode that generated or corresponds to the transaction.

(352) In the illustrated embodiment, the block entry for the transactions of body portion **306A** includes a transaction identifier that uniquely identifies the transaction, an indication of ledger changes, the identification of the channel with which the blockchain is related (channel **5**), the signatures of the endorsing peer nodes (peer node **1** and peer node **2** for the first transaction and peer node **3** for the second transaction), the signature of the ordering node that ordered the transactions (ordering node **2**), and the signature of the validating peer node (peer node **6**). As shown, given that the transactions are included in the same blockchain, the channel and validating peer node for the transactions in the body portion **2006A** is the same. However, the endorsing peer nodes are different. As described herein, this can be due to the peer nodes involved in a transaction as determined by the bytecode and/or request made by a client computing device.

(353) As described herein, the information in the block **2002A** can be used to generate one or more transaction notifications. For example, one transaction notification can include the entirety of the block **2002A**. As another example, a transaction notification can include information about the validation of the transactions in the block. For example, the transaction notification can identify the transactions of a block that are validated and/or invalidated, etc.

(354) FIGS. **21A-21D** are data flow illustrating an embodiment of a distributed ledger system **1802** processing a transaction and generating and storing a block that includes the transaction to a blockchain. In some cases, the validation process described herein with reference to FIGS. **21A-21D** can correspond to the validation of one or more transactions on a particular channel within the distributed ledger system **1802**.

(355) In the illustrated embodiment of FIG. **21A**, (1) a client computing device **2102** proposes a transaction to peer nodes **1806A** and **1806F** and receives an endorsed transaction in return. As mentioned, the peer nodes **1806A** and **1806F** can be associated with different parties or organizations. Further, the proposed transaction may relate to a proposed physical transaction between the different organizations.

(356) The peer nodes **1806A** and **1806F** process the proposed transaction and determine whether to endorse it. In certain embodiments, upon receipt of the proposed transaction, the peer nodes **1806A** can assign a transaction identifier to the proposed transaction. In certain embodiments, the client computing device **2102** can generate a transaction identifier for the proposed transaction and communicate the transaction identifier to the peer nodes **1806A** and **1806F** with the or as part of the proposed transaction. The peer nodes **1806** and ordering node **1806** can use the transaction identifier to uniquely identify the transaction throughout the validation process.

(357) In some cases, processing the proposed transaction can include executing bytecode related to the proposed transaction using one or more blocks of a respective peer node ledger **1808** or by referencing the ledger state **1904**. In certain cases, the execution of the bytecode does not modify any blocks or the ledger state **1904**, but merely verifies whether the proposed transaction could be done based on the information in the blocks and ledger state **1904**. In response to the proposed transaction, the peer nodes **1806A**, **1806F** can endorse the proposed transaction. For example, if the proposed transaction includes the proper credentials and references the correct values of the ledger

state **1904**, and identifies the proper values as part of the transaction, the peer nodes **1806A**, **1806F** can endorse the proposed transaction. As yet another example, the peer nodes **1806** can endorse a transaction based on a user determining that an entity associated with the peer node **1806** desires to proceed with the transaction. For example, if the transaction corresponds to the change in ownership, then the entities associated with the change in ownership can endorse the proposed transaction via the peer nodes **1806**.

(358) In some cases, the peer nodes **1806A**, **1806F** can endorse the proposed transaction by digitally signing the proposed transaction using a private key of a public-private key pair. In certain cases, if the peer nodes **1806A**, **1806F** do not endorse the proposed transaction (within a particular amount of time) the transaction can fail or the client computing device **2102** can resubmit the proposed transaction at a later time.

(359) In the illustrated embodiment of FIG. **21A**, the client computing device **2102** communicates with peer nodes **1806A** and **1806F**. However, it will be understood that the interactions can vary depending on the type of transaction, permissions, etc. In some cases, based on the transaction, the client computing device **2102** interacts with only one peer node **1806**. In certain embodiments, the client computing device **2102** can interact with multiple peer nodes **1806**. Further, in some embodiments, as part of the validation, one peer node **1806** can interact with another peer node **1806**. For example, if the transaction is a transfer of ownership between an entity associated with peer node **1806A** and a different entity associated with peer node **1806F**, and the transaction is initiated with peer node **1806A**, the peer node **1806A** can communicate the proposed transaction to peer node **1806F** for endorsement. In certain embodiments, an application executing on the client computing device **2102** identifies the peer nodes **1806** that are associated with a particular proposed transaction and communicates the proposed transaction to the different peer nodes **1806** for endorsement. In some cases, the peer nodes can endorse the proposed transaction in a round robin fashion. For example, after one peer node **1806** endorses the proposed transaction, it can forward the proposed transaction to another peer node for endorsement until a threshold number (e.g., all or a particular subset) of the peer nodes **1806** have endorsed the proposed transaction. In some embodiments, the ordering nodes **106** are not involved with the endorsement of the proposed transaction.

(360) With reference to FIG. **21B**, the client computing device **2102** can (2) request the ordering nodes **1806B**, **1806E** to order the transaction. As part of requesting the ordering nodes **1806** to order the transaction, the client computing device **2102** can provide the ordering nodes **1806B**, **1806E** with the endorsed transaction.

(361) Although illustrated as providing the endorsed transactions to two ordering nodes, it will be understood that the client computing device **2102** can provide the endorsed transactions to fewer or more ordering nodes **1806** as desired. In addition, in certain embodiments, one or more of the endorsing peer nodes **1806A**, **1806F** can provide the endorsed transaction to the ordering nodes **1806B**, **1806E** for ordering.

(362) The ordering nodes **1806B**, **1806E** can (3) process the endorsed transaction received from the client computing device **2102**. In some cases, processing the endorsed transaction can include ordering the endorsed transaction relative to other endorsed transactions of the distributed ledger system **1802**. For example, multiple client computing devices **2102** can be interacting with any one or any combination of the peer nodes **1806** to generate endorsed transactions. The ordering nodes **1806** can receive the endorsed transactions and order them.

(363) In certain embodiments, the ordering nodes **1806** can order the endorsed transactions based on a timestamp, such as the first, last, or an average of the timestamps of one or more of the endorsements (e.g., the timestamp associated with the peer node **1806A** and/or the peer node **1806F**), the timestamp of the proposed transaction submission or creation, etc.

(364) In addition, as part of processing the endorsed transactions, the ordering nodes **1806B**, **1806E** can generate a block for a blockchain using the endorsed transactions, including generating a

header, body, and/or other parts of the block, as discussed above. In some cases, the ordering nodes **1806B**, **1806E** can append the generated blocks to a local blockchain or ordering node ledger **1808**. In some cases, when appending the generated blocks to the local blockchain, the ordering nodes **1806** do not validate the transactions of the block. In certain embodiments, the peer nodes **1806** are not involved with the ordering of the transactions and/or the creation of the blocks from the ordered transactions.

(365) With reference to FIG. **21C**, the ordering nodes **1806B**, **1806E** (4) communicate the generated blocks to the peer nodes **1806A**, **1806C**, **1806D**, **1806F** for validation and commitment to a blockchain. As described herein, each generated block can include one or more endorsed transactions in a body portion, a header portion, and/or metadata, etc. As described herein, at least with reference to FIG. **20**, the header portion can include a hash of each transaction in the block, a hash of the hashes of each transaction, a hash of all transactions of the block or the content of a body portion of the block, a hash of a previous block of the blockchain, etc. Although both ordering nodes **1806B**, **1806E** are illustrated as providing the generated blocks to all peer nodes **1806A**, **1806C**, **1806D**, **1806F**, it will be understood that in some cases, each ordering node **1806** provides the generated blocks to a subset of the peer nodes **1806A**, **1806C**, **1806D**, **1806F** (e.g., ordering node **1806B** can send the generated blocks to peer nodes **1806A**, **1806C** and ordering node **1806E** can send the generated blocks to peer nodes **1806D**, **1806F**) or only one ordering node **1806** can provide the generated blocks to all peer nodes **1806**.

(366) The body portion can include one or more transactions or transaction data. As described herein, in some embodiments, the transaction data can include any one or any combination of: a timestamp corresponding to the transactions submission/creation, an identifier of the code (or bytecode) associated with the transaction, a signature or identification of the client computing device (or corresponding application) that initiated the transaction, a signature or identifier of the endorsing peer nodes (peer nodes **1806A**, **1806F** that signed and/or endorsed the transaction), a signature or identifier of the ordering node **1806B** that ordered the transaction and/or generated the block, a proposed change to the ledger, a channel identifier that identifies the channel associated with the blockchain, an expected input/output of the transaction, such as the content of a database of the ledger that stores the key-values associated with different transactions before and after the change is implemented, etc. Further, in some cases, the transaction data can include an identification of log data generated during bytecode execution, a bytecode response, etc.

(367) As illustrated at FIG. **21D**, the peer nodes **1806A**, **1806C**, **1806D**, **1806F** can (5) validate the transactions in the block and append or commit the block to a peer node ledger **1808** and/or a peer node blockchain. In certain embodiments, the peer nodes **1806A**, **1806C**, **1806D**, **1806F** can validate the transactions by comparing the expected inputs (e.g., value indicated in the transaction for a particular key of the ledger state **1904** compared to the actual value of the key in the ledger state **1904**). In some cases, if the value or state of the key in the ledger state **1904** matches the value or state identified by the transaction, the peer node **1806** can validate the transaction. In certain cases, the peer nodes can validate the transactions based on permissions or other information associated with the endorsing peer nodes **1806A**, **1806F**, etc.

(368) In addition, in some cases, the peer nodes **1806A**, **1806C**, **1806D**, **1806F** can update the ledger state **1904** based on the transactions. For example, as described herein, the ledger state **1904** can store key-values corresponding to the subject of one or more transactions. When a transaction affects a particular key-value pair, the peer nodes **1806** can update the key-value pair in the respective ledger state **1904** and append the corresponding block to the blockchain of the respective peer node ledger **1808**. As described herein, the ledger state **1904** can reflect the current state or value of a key based on the combination of valid transactions in a blockchain.

(369) Throughout the validation process, the nodes **1806** can generate different types of data, such as, but not limited to transaction notifications, log data, and/or metrics data.

(370) In some cases, the peer nodes **1806** can generate one or more transaction notifications. The

transaction notifications can correspond to individual transactions of a block, the entire block, or parts of a transactions, such as the bytecode used as part of a transaction, etc. In some cases, the transaction notifications can include the entire content of a block (e.g., the header portion, body portion, transactions, metadata, etc.), or a summary of information, such as an indication of which transactions were validated and/or posted to a peer node blockchain. In certain embodiments, the notifications can be stored in a pub-sub or buffer and/or the peer nodes **1806** can notify the client computing device **2102** based on the generated transaction notifications, and provide client computing device **2102** with information about the transaction as part of the block of a blockchain. In some cases, the peer node **1806** can indicate to the client computing device **2102** whether the transaction was validated or invalidated, etc.

(371) In addition to generating notifications, the nodes **1806** can generate log data. The log data can correspond to or identify different transactions that are being processed by the nodes **1806** or other activities related to the node, such as errors, etc. For example, the log data generated by a peer node **1806** can indicate what the peer node **1806** doing for a particular proposed transaction (e.g., receive transaction, assign transaction identifier, endorse transaction, validate/invalidate transaction, post block with transaction to blockchain, read/write proposed changes of the transaction to the ledger state **1904**, etc.). Similarly, the ordering nodes **1806** can generate log data indicative of activities it is executing relative to the transactions (e.g., receive endorsed transaction, order transaction, generate block, add transaction to a block, communicate transaction to peer nodes as part of the block, post transaction to blockchain as part of a block, etc.). Though log data can capture the activity of a node as the node processes transactions, the log data for the node can, in some cases, only capture the activity of the one node. Depending on the implementation of the nodes **1806**, the log data can be stored in a data store of the nodes, and/or converted and stored as part of log data of an isolated execution environment system, etc.

(372) Moreover, as the nodes **1806** process data, they can generate certain metrics. For example, the nodes **1806** can generate CPU usage, diskspace, etc. and other metrics. Though the metrics for a node result from processing performed by the node, metrics data may not capture any information about transactions that were processed. In some cases, the metrics are stored in a data store of the nodes **1806**.

(373) The data intake and query system **108** can ingest and correlate the data generated by a distributed ledger system **1802**. In some cases, the data intake and query system **108** can ingest the data using different components. For example, the data intake and query system **108** can use a monitor to ingest one type of data and use a forwarder, connector, and/or data adapter for other types of data.

(374) Based on the collected data, the data intake and query system **108** can identify correlations between transactions that are included in a blockchain and corresponding log data and metrics data. This information can provide insight into the inner workings of the distributed ledger system **1802**, identify performance issues, security issues, errors, etc. By identifying faults, errors, and issues with the different components of the distributed ledger system **1802**, the data intake and query system **108** can improve the distributed ledger system **1802** as a whole. For example, based on the identified issues, system configurations can be adjusted, components can be fixed or reconfigured, etc. In this way, the data intake and query system **108** can improve the speed, efficiency, throughput, and processing power of the distributed ledger system **1802**. In addition, by correlating the different data types or associating data from different nodes of the distributed ledger system **1802**, the data intake and query system **108** can track the throughput of the system, identify bottlenecks, and be used to make adjustments to the distributed ledger system **1802**.

(375) FIG. 22 is a data flow diagram illustrating data ingestion from the distributed ledger system **1802** by the data intake and query system **108** operating in accordance with aspects of the present disclosure. As noted herein above, the data intake and query system **108** may implement a Getting-Data-In (GDI) component (such as data adapter, monitor, forwarder, connector, or the like) in order

to ingest the distributed ledger transaction data, e.g., by reading log files maintained by one or more nodes of the distributed ledger, listening to the blocks, transactions, and events that are broadcasted to all participating nodes of a distributed ledger, and/or performing other actions. The ingested raw data may be aggregated, decoded, visualized, and/or further processed by the data intake and query system.

(376) In the illustrated embodiment of FIG. 22, the node **1806** is shown as a peer node **1806** with a peer node ledger **1808**, a ledger state **1904**, and a buffer **2204**. The node **1806** can generate different types of data, including transactions, blocks, metrics data, and/or log data. In an illustrative example, a transaction can include the transaction identifier, a timestamp, the source account identifier, the destination account identifier, and a bytecode invoking a smart contract by specifying a function name and parameter values.

(377) The log data can include information generated by the node as it processes requests, transactions, etc. The log data can, for example, include information about errors, or other issues. In some cases, the log data can include a transaction identifier indicating a particular transaction associated with the generated log data. For example, the log data can indicate that a particular transaction associated with a particular transaction identifier was received, rejected, forwarded, processed, endorsed, ordered, included in a block, validated, invalidated, pruned, caused an error, rejected, used to edit the ledger state **1904**, etc. The log data can also include information about other occurrences within the node **1806**, such as, but not limited to, interactions with other nodes **1806**, setup, administrative communications, configuration settings/changes, etc. As described herein, in some embodiments, the log data may be unstructured raw machine data, whereas the transaction notifications may be structured.

(378) As described herein, the metrics can include information about the performance metrics of the node **1806** and/or the distributed ledger system **1802**, such as, but not limited to, CPU-related performance metrics; disk-related performance metrics; memory-related performance metrics; network-related performance metrics; energy-usage statistics; data-traffic-related performance metrics; overall system availability performance metrics; cluster-related performance metrics; and virtual machine performance statistics, etc.

(379) The different types of data generated by the node **1806** and/or distributed ledger system **1802** can be accessible via different paths or stored in different locations of the node **1806**. For example, the data can be located in a data store, pub/sub, buffer, or real-time data stream. In some embodiments, such as when the distributed ledger system **1802** is implemented in an isolated execution environment system, the data can be wrapped or converted to another format, such as JSON, and stored as a JSON (or other type) file. In some such cases, a data adapter **2202**, connector, or monitor can be used to extract the data generated by the node from the wrapper. In some such embodiments where the distributed ledger system **1802** is implemented using Kubernetes or Docker, log data generated by the node **1806** can be wrapped in a JSON wrapper and stored as a Docker or Kubernetes log file, which may, for example, be accessible through an API. In some such case, a data adapter or monitor can use the API or another mechanism to extract the log data generated by the node from the JSON wrapper and/or Docker or Kubernetes log file.

(380) As described herein, in some cases, the data obtained from the node **1806** can be available via a messaging buffer **2204**. In certain embodiments, the buffer **2204** operates according to a publish-subscribe (“pub-sub”) messaging system. For example, a channel may be represented as one or more “topics” within a pub-sub system, and new transaction notifications may be represented as a “message” within the pub-sub system. The distributed ledger system monitor **1804** may subscribe to a topic representing desired information (e.g., a particular channel, all transaction notifications, etc.) to receive messages within the topic. Thus, the distributed ledger system monitor **1804** can be notified of new data categorized under the topic within the buffer **2204**. A variety of implementations of the pub-sub messaging system may be usable within the buffer **2204**. As will be appreciated, use of a pub-sub messaging system can provide many benefits, including the ability to



retrieve data quickly from the node **1806** while maintaining or increasing data resiliency. In some embodiments, the distributed ledger system monitor **1804** can provide the data to the data intake and query system **108** through a module that provides an intake API to the data intake and query system **108**.

(381) In certain embodiments, the data can be collected from the node **1806** by installing one or more forwarders **1904** on the node **1806** and/or or using an HTTP event collector (indicated by arrow **2206**). For example, the metrics or log data can be obtained from the node **1806** using a forwarder **1904** or HTTP event collector.

(382) As described herein, the data obtained from the node **1806** can be stored in one or more buckets of the data intake and query system **108**. In some cases, the log data can be stored in one set of buckets associated with one index, the metrics data can be stored in a second set of buckets associated with a second index, and the transaction notifications can be stored in a third set of buckets associated with a third index. However, it will be understood that the data obtained from the node **1806** can be stored in a variety of ways and formats.

(383) Moreover, the data intake and query system **108** can populate one or more inverted indexes based on the received data. In some embodiments, as the data intake and query system **108** ingests the data from the node **1806**, it can extract a transaction identifier using one or more regex rules. For example, the data intake and query system **108** can use one or more regex rules to extract a transaction identifier from log data and/or transaction notifications. As the log data and transaction notifications are different data types or have a different sourcetype, the data intake and query system **108** can use different regex rules to extract the corresponding transaction identifier (e.g., use one regex rule to extract transaction identifier from log data and a different regex rule to extract transaction identifiers from transaction notifications). In some cases, the distributed ledger system monitor **1804** can extract the transaction identifiers from the transaction notifications. The data intake and query system **108** can include the extracted transaction identifiers as keywords or field-value pairs in one or more inverted indexes, such as the inverted index describe herein with reference to FIG. 5B. In certain embodiments, the data intake and query system **108** can extract a node identifier for each node **1806** of the distributed ledger system **1802**. The extracted node identifier can also be stored in one or more inverted indexes. Similarly, the data intake and query system **108** can extract other data from the transaction notifications and/or log data (non-limiting examples: endorsing, ordering, validating node identifiers, channel identifiers, ledger state edit time **1904**, etc.) and store the extracted data in one or more inverted indexes.

(384) As described herein, the data intake and query system **108** can correlate different types of data of a particular node or across different nodes and/or associate the same type of data of a particular node or across different nodes. For example, the data intake and query system **108** can correlate log data with transaction notifications and/or metrics data from the same node or from different nodes. Similarly, the data intake and query system **108** can associate log data from different nodes or transaction notifications from different nodes, etc.

(385) In some cases, a node **1806** can generate multiple log data entries for each transaction notification. For example, in some embodiments, a transaction notification can correspond to the commitment of a block to a blockchain, whereas the log data can correspond to one or more processing tasks or other activities performed by the node **1806**. As a node **1806** can perform multiple processing tasks or activities before committing a block to a blockchain, there can be multiple of entries in log data or multiple log data events for each transaction notification or transaction notification event. In addition, as each peer node **1806** can maintain its own blockchain, each peer node **1806** can generate a transaction notification that identifies the same transaction (or includes the same transaction identifier). Accordingly, log data or log data events of one node can be correlated with multiple transaction notifications from different nodes.

(386) In some embodiments, when correlating log data and transaction notifications of a particular node, multiple entries of log data can be correlated with one (or more, depending on the

embodiment) transaction notification, and when correlating log data and transaction notifications across multiple nodes, one entry of log data (or one log data event) can be correlated with multiple transaction notifications (or transaction notification events) from different nodes. In some such embodiments, the data intake and query system **108** can correlate log data and transaction notifications of each node before correlating log data and transaction notifications across nodes. In a similar manner, multiple sets of metrics data of a particular node can be associated with a particular entry in log data (or log data event) or transaction notification of the particular node.

(387) As described herein, correlating the data obtained from the nodes **1806** can provide significant insights and improvements for the distributed ledger system **1802**. In some cases, the data obtained from a single node can be correlated to provide node diagnostics, identify the structure or architecture of the node **1806** and/or parts of the distributed ledger system **1802**, identify node failures or bottlenecks, recreate or rebuild the blockchain or ledger state **1904**, determine the history of a transaction with reference to the node **1806** or partial history, etc.

(388) In some cases, to correlate transaction notifications with log data, the data intake and query system **108** can identify events that are associated with the same transaction identifier. As described herein, some of the events can correspond to log data (also referred to herein as log data events) from a node **1806** and other events can correspond to transaction notifications of a node **1806** (also referred to herein as transaction notification events). Based on a determination that the log data events and the transaction notification events include the same transaction identifier, the data intake and query system **108** can correlate the different events.

(389) In some cases, the data intake and query system **108** can correlate the metrics with the log data and/or transaction identifiers based on one or more timestamps. For example, as described herein, events can include data associated with a timestamp and metrics can be stored in association with a timestamp. Accordingly, the data intake and query system **108** can correlate the metrics with the log data and/or transaction notifications using the corresponding timestamps. In this way, the data intake and query system **108** can determine the relevant metrics of the node **1806** at the time particular log data and/or transaction notification was generated. This correlation can provide insights into the state of the node **1806** as when the log data and/or transaction notification was generated.

(390) Further, the correlation of different data can provide different insights into the state of the distributed ledger system **1802**, a transaction, and/or a node **1806**. For example, by correlating the log data events and transaction notification events of a node, the data intake and query system **108** can identify node failures in relation to particular transactions, node throughput, determine the amount of validated vs. invalidated nodes, the timing/frequency of the generation of a block or commitment of the block to a blockchain.

(391) In certain embodiments, correlating transaction notifications and log data can provide insights into the content of a block of a blockchain. For example, the content of the blockchain may be encoded, encrypted, or otherwise obfuscated for privacy or security purposes. By correlating a transaction of a block with corresponding log data, the data intake and query system **108** can determine some or all of the content of the transaction of the block. For example, the log data may include information about the transaction that was encrypted or otherwise obfuscated in the block.

(392) Similarly, the correlation or association of data across nodes can provide insights into the state of the distributed ledger system **1802**, a transaction, a transaction history, etc. In some cases, the data intake and query system **108** can associate the same type of data across multiple nodes **1806** of the distributed ledger system **1802**. For example, the data intake and query system **108** can associate log data events from a first peer node **1806** with log data events from a second peer node **1806**. In some cases, the association of the same type of data can be used to identify the history of a transaction as it is received, endorsed, ordered, validated, included in a block, and/or committed to a blockchain.

(393) Associating transaction notification events across different nodes can provide insights into

the functioning of the distributed ledger system **1802**. In some cases, the association can be used to identify errors in a particular node **1806**. For example, if all but one peer node **1806** of a distributed ledger system **1802** have committed a particular block to a respective blockchain, the data intake and query system **108** can identify a potential fault or error with the particular peer node. Similarly, the correlation or association of data across nodes **1806** can enable the data intake and query system **108** to compare the throughput and processing of each node **1806** to identify slower/faster nodes, bottlenecks, etc. Accordingly, by obtaining the different types of data from the nodes of a distributed ledger system **1802** and correlating the data, the distributed ledger system **1802** can be improved. Specifically, the correlation can identify vulnerabilities, faults, errors, etc., of the distributed ledger system **1802**. Correlating the data across the nodes **1806** can also enable the identification of potential security issues, such as, but not limited to, validated transactions that were not endorsed, digital certificate or signature abnormalities, an abnormal volume of transactions, significant transactions or interactions with computers from a particular geographic area or block of IP addresses, etc.

(394) In addition, by correlating the data across the nodes **1806**, the data intake and query system **108** can determine the architecture of the distributed ledger system **1802**. For example, the data intake and query system **108** can identify the different peer nodes **1806** and ordering nodes **1806** of the distributed ledger system **1802**, etc. For example, the data intake and query system **108** can identify log data events and transaction notification events that are associated with the same channel (e.g., based on parsing relevant events and/or using one or more inverted indices). The data intake and query system **108** can then identify different peer nodes **1806** and ordering nodes **1806** associated with the identified log data events and transaction notification events. Based on this information, the data intake and query system **108** can identify the different nodes **1806** of a channel. Further, by doing this for some or all channels, the data intake and query system **108** can identify the nodes **1806** of the distributed ledger system **1802**. In addition, the data intake and query system **108** can determine with which channels each node **1806** is associated, which nodes **1806** share bytecodes, frequent transactions between nodes, the size and/or number of nodes **1806** involved in the individual transactions, etc.

(395) By associating and correlating the events, the data intake and query system **108** can identify the components of the distributed ledger system **1802** (and/or the status of the components). For example, by associating and correlating events, the data intake and query system **108** can determine that a particular distributed ledger system **1802** includes three ordering nodes **1806**, ten peer nodes **1806**, four channels, etc. Similarly, the data intake and query system **108** can determine the status for the different components. For example, the data intake and query system **108** can determine the number of errors, warnings, responsiveness, etc. of one or more the three ordering nodes **1806** and ten peer nodes **1806**.

(396) In certain embodiments, the data intake and query system **108** can obtain and correlate additional types of data. For example, as described herein, blocks of a blockchain can include one or more digital signatures of a peer node **1806** and/or an ordering node **1806**. The data intake and query system **108** can use the digital signature to identify a Certificate Authority associated with the digital signature (e.g., a Certificate Authority that issued a digital certificate to the peer node **1806** and/or ordering node **1806** to which the digital signature corresponds). One or more components of the data intake and query system **108** (e.g., distributed ledger system monitor **1804**, forwarder, and/or indexer) can query the Certificate Authority to obtain additional identifying attributes of the signer such as name, address, email, company name, phone number, title, etc. This information can be stored by the data intake and query system **108** and/or correlated with the other data obtained from the distributed ledger system **1802** to diagnose issues with specific transactions and/or answer business analytics questions.

(397) In some cases, the data intake and query system **108** can correlate the data to identify relationships between the components of the distributed ledger system **1802**, and generate a

visualization based on the relationships. For example, the data intake and query system **108** can determine which nodes are associated with which channels and can therefore communicate with each other with respect to a particular blockchain. In some cases, the data intake and query system **108** can store the determined relationships and/or other attributes of the different components in a table. In certain cases, the table can be used to generate a visualization. In certain embodiments the visualization can indicate the relationships of the components of the distributed ledger system **1802**. For example, the visualization can indicate peer nodes **1806** and ordering nodes **1808** that are part of a shared channel or consortium. In addition, the various processing steps of a transaction can be tracked across the different nodes of the distributed ledger system **1802** and visualized. In this way, the data intake and query system **108** can identify issues and errors with the distributed ledger system **1802**, etc.

(398) In some embodiments, the visualization can indicate the overall health of individual components of the distributed ledger system **1802**. For example, the visual representation of the components can be colored green to indicate “healthy” (e.g., fewer than a threshold number of errors/warnings or no errors/warnings) and red to indicate “unhealthy” (e.g., greater than a threshold number of errors, warnings, etc.). The visualization can also enable the user to drill down on visual representations of the components of the distributed ledger system **1802** to display information, metrics, and log data to determine the performance of the component and/or troubleshoot problem areas. In certain embodiments, such as when a user “drills down” to a particular component, the visualization can display log data, metrics data, trends of log data or metrics data, etc.

(399) As noted herein above, the distributed ledger system **1802** may support a special account type, which is referred to as “contract account.” A message to a contract account activates its executable code implementing a “smart contract,” which may evaluate specified conditions and perform various actions (e.g., transfer cryptocurrency tokens between accounts, write data to internal storage, mint new cryptocurrency tokens, perform calculation, create new smart contracts, etc.). The nodes of the distributed ledger system **1802** may collectively implement a distributed virtual machine (e.g., the EVM) for executing the code implementing smart contracts. A smart contract can be created in a high level programming language (such as Solidity) and then compiled into the EVM bytecode.

(400) FIG. **23** is a data flow diagram illustrating transaction decoding by a distributed ledger connector of the data intake and query system **108** operating in accordance with aspects of the present disclosure. In an illustrative example, the distributed ledger connector **2302** may receive, from a distributed ledger node **2304**, a transaction **2310**. In another illustrative example, the transaction **2310** may be retrieved from a log file (not shown) maintained by the distributed ledger node **2304**.

(401) The transaction **2310** may include the transaction nonce **2312**, the transaction processing fee **2314**, the destination account identifier **2316**, the transaction value **2318**, the transaction data **2320**, and the cryptographic signature **2322** of the transaction originating node. The transaction nonce **2312** may specify a sequence number of transactions sent from the given source address. The transaction processing fee **2314** may specify the maximum amount of crypto currency which the transaction originator is willing to pay for processing the transaction. The destination account identifier **2316** may identify an external account or a contract account. The transaction value **2318** may specify an amount of crypto currency to be transferred to the destination account. The transaction data **2320** may be empty for payment transactions; for smart contract transactions, the transaction data **2320** contains a message invoking a smart contract.

(402) In order to decode a smart contract transaction, the distributed ledger connector **2302** may retrieve the bytecode **2324** implementing the smart contract associated with the distributed ledger account **2316** identified by the transaction as the destination account. In an illustrative example, the bytecode may be retrieved by performing a JSON-RPC call to the destination node specified by the

transaction **2310**. In another illustrative example, the bytecode may be retrieved from a log file maintained by the destination node or another node of the distributed ledger system. Upon retrieving the bytecode, the distributed ledger connector **2302** may compute its digital fingerprint **2326**.

(403) FIG. **24** is a data flow diagram of computing a digital fingerprint of a smart contract, in accordance with aspects of the present disclosure. As schematically illustrated by FIG. **24**, a digital fingerprint of a smart contract can be represented by a cryptographic hash of all distributed ledger function signatures and distributed ledger event signatures contained in the bytecode **2400** implementing the smart contract. Accordingly, the distributed ledger connector **2302** may parse the bytecode **2400** and extract the distributed ledger function signatures **2410A-2410P** and distributed ledger event signatures **2420A-2420Q** (collectively referred to as Function and Even Signatures **2450**).

(404) In an illustrative example, a function signature may be represented by the function name followed by a parenthesized list of parameter types that are split by a predefined delimiter (e.g., a comma):

(405) Function Name (param\_1\_type, param\_2\_type, . . . param\_N\_type), where N is the number of parameters.

(406) Similarly, a distributed ledger event signature may be represented by the distributed ledger event topic followed by a parenthesized list of parameter types that are split by a predefined delimiter (e.g., a comma):

(407) EventTopic (param\_1\_type, param\_2\_type, . . . param\_M\_type), where M is the number of parameters.

(408) Alternatively, various other specifications of distributed ledger function signatures and/or distributed ledger event signatures may be utilized, provided that a signature unambiguously reflects at least the function name (or the distributed ledger event name) and parameter types.

(409) As noted herein above, the digital fingerprint of a smart contract can be represented by a cryptographic hash of all distributed ledger function signatures and distributed ledger event signatures. A cryptographic hash may be represented by an irreversible function mapping its argument represented by a bit string of arbitrary size to a hash value represented by a bit string of a pre-determined size, such that two different arguments are very unlikely to produce the same hash value. In an illustrative example, a Secure Hash Algorithm (SHA) function, such as SHA-3, may be utilized for computing the digital fingerprint. Alternatively, various other cryptographic hash functions may be utilized.

(410) Upon extracting, from the bytecode **2400**, the distributed ledger function signatures and distributed ledger event signatures **2430**, the distributed ledger connector **2302** may compute the value of the chosen hash function **2440** of the character string produced by concatenating all extracted distributed ledger function signatures and distributed ledger event signatures, thus producing the digital fingerprint **2326**:

(411) Fingerprint=H (concatenation (signature\_1, signature\_2, . . . , signature\_K), where H is the chosen cryptographic hash, and signature\_1, signature\_2, . . . , signature\_K denote the distributed ledger function signatures and distributed ledger event signatures extracted from the contract bytecode.

(412) Referring again to FIG. **23**, the distributed ledger connector **2302** may utilize the computed digital fingerprint **2326** for associating the smart contract with a known application binary interface (ABI) definition. The data intake and query system **108** may maintain a local database **2330** of ABI definitions of smart contracts and/or access a remote database (not shown in FIG. **23**) of ABI definitions of smart contracts. In various illustrative examples, the remote database may be a publicly accessible database or a private database maintained by a customer utilizing the data intake and query system **108** or by a third party.

(413) The ABI definition of a smart contract may be represented by a human-readable textual

representation (e.g., a JSON file) describing the smart contract and its functions. For each function, its name, parameters types and values, and other pertinent information is specified. Accordingly, for every smart contract having its ABI definition **2332** stored in the database **2330**, the data intake and query system may compute the digital fingerprint **2334** following the procedure described herein above with reference to FIG. **24**. The computed digital fingerprints **2334** may be stored in the database **2330** in association with the respective ABI definitions **2332**, such that the digital fingerprint **2334A** is associated with the ABI definition **2332A**, the digital fingerprint **2334B** is associated with the ABI definition **2332B**, etc. The database **2330** may be indexed by the values of the digital fingerprints **2334** in order to facilitate efficient identification and retrieval of a smart contract ABI definition associated with a specified digital fingerprint.

(414) Thus, the distributed ledger connector **2302** may search the database **2330** for a digital fingerprint that matches the computed digital fingerprint **2450** of the bytecode **2322** implementing the smart contract associated with the distributed ledger account **2316**. Upon identifying the matching digital fingerprint **2334R**, the distributed ledger connector **2302** may retrieve the associated ABI contract definition **2332R**. The distributed ledger connector **2302** may extract, from the ABI contract definition **2332R**, function signatures specifying, for each function exposed by the ABI definition, its name as well as the names and types of its parameters. The extracted information may be utilized for decoding the transaction data **2320** thus producing decoded transaction data **2350**.

(415) As noted herein above, for smart contract transactions, the transaction data **2320** contains a message invoking a smart contract. Accordingly, the distributed ledger connector **2302** may, upon identifying within the transaction data **2320**, a smart contract invocation, extract the function signature and parameter values associated with the identified transaction data. In an illustrative example, a predefined number of bytes (e.g., four bytes) of the function call data referenced by the transaction data specify the hash of the signature of the function to be called. Following the function signature (e.g., starting from the fifth byte), the function call contains encoded parameter values.

(416) The distributed ledger connector **2302** may identify, within the ABI definition **2332R** of the identified matching smart contract, a function signature matching the function signature extracted from the transaction data **2320**. The definition of the identified matching function (e.g., parameter names and types) may be utilized by the distributed ledger connector **2302** for decoding the transaction data **2320**. In particular, the distributed ledger connector **2302** may decode the parameter values extracted from the transaction data **2320** in accordance with the parameter types specified by the ABI definition **2332R** for the identified matching function. The distributed ledger connector **2302** may further associate each parameter value with a corresponding parameter name specified by the ABI definition **2332R** for the identified matching function. The decoded transaction data, including the function name, the parameters names, types, and values, may be fed to a data intake and query system **108** for visualization and/or further processing.

(417) FIG. **25** is a flow diagram of an embodiment of a method **2500** of decoding distributed ledger transactions, in accordance with aspects of the present disclosure. Method **2500** and/or each of its individual functions, routines, subroutines, or operations may be performed by one or more processors of a computing device the distributed ledger connector **2302** of the data intake and query system **108**. In certain implementations, method **2500** may be performed by a single processing thread. Alternatively, method **2500** may be performed by two or more processing threads, each thread executing one or more individual functions, routines, subroutines, or operations of the method. In an illustrative example, the processing threads implementing method **2500** may be synchronized (e.g., using semaphores, critical sections, and/or other thread synchronization mechanisms). Alternatively, the processing threads implementing method **2500** may be executed asynchronously with respect to each other. Therefore, while FIG. **25** and the associated description lists the operations of method **2500** in certain order, various implementations

of the method may perform at least some of the described operations in parallel and/or in arbitrary selected orders.

(418) Although described as being implemented by the distributed ledger connector **2302** of the data intake and query system **108**, it will be understood that one or more elements outlined for method **2500** can be implemented by one or more computing devices/components that are associated with a data intake and query system **108**, such as the search head **210**, indexer **206**, etc. Thus, the following illustrative embodiment should not be construed as limiting.

(419) At block **2510**, the computing device implementing the distributed ledger connector **2302** receives distributed ledger transaction. The transaction may include the transaction nonce, the transaction processing fee, the destination account identifier, the transaction value, the transaction data, and the cryptographic signature of the transaction originating node. The transaction nonce may specify a sequence number of transactions sent from the given source address. The transaction processing fee may specify the maximum amount of crypto currency which the transaction originator is willing to pay for processing the transaction. The destination account identifier may identify an external account or a contract account. The transaction value may specify an amount of crypto currency to be transferred to the destination account. The transaction data may be empty for payment transactions; for smart contract transactions, the transaction data contains a message invoking a smart contract.

(420) At block **2520**, the computing device receives the bytecode module associated with the distributed ledger account identified by the transaction. In an illustrative example, the bytecode may be retrieved by performing a JSON-RPC call to the destination node specified by the distributed ledger transaction. In another illustrative example, the bytecode may be retrieved from a log file maintained by the destination node or another node of the distributed ledger system.

(421) At block **2530**, the computing device computes the digital fingerprint of the retrieved bytecode. The digital fingerprint may be represented by a cryptographic hash of all distributed ledger function signatures and distributed ledger event signatures contained by the bytecode. Thus, upon extracting the distributed ledger function signatures and distributed ledger event signatures from the bytecode, the computing device may compute the value of a chosen hash function of the character string produced by concatenating all extracted distributed ledger function signatures and distributed ledger event signatures, thus producing the digital fingerprint, as described in more detail herein above.

(422) At block **2540**, the computing device identifies, among a plurality of ABI definitions stored in the ABI definition database, an ABI definition having the ABI digital fingerprint that matches the computed bytecode digital fingerprint, as described in more detail herein above.

(423) At block **2550**, the computing device produces decoded transaction data by decoding, using the identified ABI definition, the transaction data. In an illustrative example, the computing device may retrieve, from the ABI definition database, the ABI contract definition associated with the identified digital fingerprint matching the computed bytecode fingerprint. The computing device may then extract, from the retrieved ABI contract definition, the function signatures specifying, for each function exposed by the ABI definition, its name as well as the names and types of its parameters. The extracted information may be utilized for decoding the transaction data, as explained in more detail herein below with reference to FIG. **26**

(424) FIG. **26** is a flow diagram of an embodiment of a method **2600** of decoding transaction data, in accordance with aspects of the present disclosure. Method **2600** and/or each of its individual functions, routines, subroutines, or operations may be performed by one or more processors of a computing device the distributed ledger connector **2302** of the data intake and query system **108**. In certain implementations, method **2600** may be performed by a single processing thread. Alternatively, method **2600** may be performed by two or more processing threads, each thread executing one or more individual functions, routines, subroutines, or operations of the method. In an illustrative example, the processing threads implementing method **2600** may be synchronized

(e.g., using semaphores, critical sections, and/or other thread synchronization mechanisms). Alternatively, the processing threads implementing method **2600** may be executed asynchronously with respect to each other. Therefore, while FIG. **26** and the associated description lists the operations of method **2600** in certain order, various implementations of the method may perform at least some of the described operations in parallel and/or in arbitrary selected orders.

(425) Although described as being implemented by the distributed ledger connector **2302** of the data intake and query system **108**, it will be understood that one or more elements outlined for method **2600** can be implemented by one or more computing devices/components that are associated with a data intake and query system **108**, such as the search head **210**, indexer **206**, etc. Thus, the following illustrative embodiment should not be construed as limiting.

(426) At block **2610**, the computing device implementing the distributed ledger connector parses the transaction data to identify a smart contract function invocation (encoded signature hash and parameters).

(427) At block **2620**, the computing device extracts the signature and parameter values of a function invoked by the transaction data. In an illustrative example, a predefined number of bytes (e.g., four bytes) of the identified function call specify the hash of the signature of the function to be called. Following the function signature (e.g., starting from the fifth byte), the function call contains encoded parameter values.

(428) At block **2630**, the computing device identifies, in the ABI definition having the ABI digital fingerprint that matches the bytecode digital fingerprint, a function signature matching the signature extracted from the transaction data.

(429) At block **2640**, the computing device utilizes the definition of the identified matching function (e.g., parameter names and types) for decoding the transaction data. In particular, the computing device may decode the parameter values extracted from the transaction data in accordance with the parameter types specified by the ABI definition for the identified matching function. The computing may further associate each parameter value with a corresponding parameter name specified by the ABI definition for the identified matching function. The decoded transaction data, including the function name, the parameters names, types, and values, may be fed to a data intake and query system for visualization and/or further processing.

(430) In some cases, the data intake and query system **108** can generate one or more visualizations of the results. In certain cases, the visualizations can indicate the path or history of the transaction through the distributed ledger system **1802**, the architecture of the distributed ledger system **1802**, and/or the status of individual nodes **1806** and/or the distributed ledger system **1802** as a whole. In certain embodiments, the visualization can include a display object for each node of the distributed ledger system **1802** with one more indicators that indicate a status of the node, such as the number of errors or faults at the node, the number of transactions processed or being processed, the number of channels associated with each node, etc.

(431) The data intake and query system **108** can use the events to identify errors, bottlenecks or other issues in the distributed ledger system **1802**. For example, the data intake and query system **108** can identify nodes with smaller or greater throughput, nodes with the most errors, etc. In some cases, the data intake and query system **108** can use the events to track the lifecycle of a transaction, including the initial submission of transaction to a node **1806**, endorsement of the transaction, ordering of the transaction, validation of the transaction, and inclusion of the transaction into the ledger **1808**. If the transaction stops or slows down unacceptably at any point in the journey, the data intake and query system **108** can diagnose the reason for the slow down and generate an alert. Potential issues may include, errors in the bytecode execution, latency with querying the ledger state **1904**, network latency, resource contention in the underlying distributed ledger system **1802**, authentication/authorization issues, etc.

(432) In addition, the results can include an identification of errors associated with the processing of the transaction, errors with or created by the bytecode, errors in the bytecode execution,



throughput of the node, time taken to process the transaction at different times, latency with querying the ledger state **1904**, etc. In some cases, based on the association, the data intake and query system **108** can determine the architecture (or a portion thereof) of the distributed ledger system **1802**. For example, the information from the various events can indicate which peer nodes endorsed the transaction and are therefore part of the distributed ledger system **1802**, which ordering node ordered the transaction and is therefore part of the distributed ledger system **1802**. In addition, the data intake and query system **108** can identify the number and identity of various channels with which the node **1806** is associated.

(433) In certain embodiments, the data intake and query system **108** can use the events to determine a type of node processing a particular transaction (e.g., peer node and/or ordering node). For example, if a node does not have any transaction notification events associated with it, the data intake and query system **108** can determine that it is an ordering node, or if it does have transaction notification events associated with it, the data intake and query system **108** can determine that it is a peer node, etc. In some cases, by associating and/or correlating the events, the data intake and query system **108** can recreate the blockchain.

#### Data Ingestion Overview

(434) As previously mentioned, the data intake and query system **108** can ingest data, process the data (e.g., perform various transformations or manipulations of the data), and output the data. For example, the data intake and query system **108** can ingest data shared across multiple customers (e.g., public data, semi-public data, shared data, etc. associated with a public data source, a semi-public data source, a shared data source, etc.). Specifically, the data may be generated by a distributed ledger system or a computing device (e.g., one or more computing devices that host, store, generate, or are otherwise associated with data that is not stored on a blockchain, offchain data, etc.). In some cases, the data intake and query system **108** can ingest the data generated by the distributed ledger system or the computing device using different components. For example, the data intake and query system **108** can ingest the data generated by the distributed ledger system or the computing device using a monitor, a forwarder, a connector, a data adapter, etc. In some cases, different types of components can be used to ingest different types of data. For example, the data intake and query system **108** can use a monitor to ingest a first type of data, a forwarder to ingest a second type of data, etc.

(435) The data ingested and processed by the data intake and query system **108** may be accessible for the data intake and query system **108** for use in executing queries. For example, the data intake and query system **108** may receive an input from a user computing device and generate a query based on the received input. The data intake and query system **108** may execute the defined query on the ingested and processed data. Therefore, the data intake and query system **108** can ingest data from distributed ledger systems and execute defined queries on the ingested data.

(436) A system may be limited to ingesting the data in response to a request to ingest the data (e.g., for execution of a query). For example, the system may be limited to receiving a request for execution of a query associated with data from a distributed ledger system, ingesting the data from the distributed ledger system in response to the request, and executing the query. However, such a system may prove unsatisfactory as the execution of the query may be delayed based on the ingestion of the data from the distributed ledger system as the system may not ingest the data until a request for execution of a query is received. As the system may not ingest the data until the request for execution of a query is received, the data from the distributed ledger system may not be indexed and the execution of the query may be inefficient.

(437) Each user requesting a query of data from a distributed ledger system may individually and separately initiate ingestion of the data. As hundreds or thousands of users may execute queries on data from a distributed ledger system, the individual and separate ingestion of the data from the distributed ledger system may be inefficient. Such an ingestion of the data from the distributed ledger system may be inefficient.

(438) A system may be limited to ingesting data from a single distributed ledger system. While the distributed ledger system may be a public data source (e.g., a public blockchain), the system may be unable to execute a query across data from multiple distributed ledger systems as the data from multiple distributed ledger systems may have different data formats, different data, different query parameters, etc. For example, a system may be limited to ingesting data from a single distributed ledger (e.g., Ethereum). The system may include a particular component (e.g., a connector) configured for ingestion of data from the single distributed ledger. However, such a system may prove unsatisfactory when a system is attempting to query data corresponding to multiple distributed ledger systems. For example, a query may include a request to correlate, contrast, compare, and/or combine data from multiple distributed ledger systems. For example, the system may not be able to execute queries to correlate data from a first distributed ledger system and data from a second distributed ledger system. Specifically, the system may not be able to identify correlations between data from unrelated distributed ledger systems (e.g., distributed ledger systems utilizing different technologies such as Ethereum and Bitcoin). Therefore, in such a system, it may not be possible for the system to ingest data from multiple distributed ledger systems and enrich the data.

(439) Such a system may be limited to ingesting the data according to a single processing schema. For example, a system may receive a query and initiate a component according to a single processing schema (identifying a set of processing parameters) to ingest the data from the distributed ledger system. The processing schema may be based on the query received by the system. For example, the processing schema may cause the component to ingest data according to one or more parameters of the query. While such a single processing schema may enable execution of a single query, such a processing schema may prove inefficient where a system or systems are executing multiple queries on data ingested from the same distributed ledger system as each execution of a query may result in a separate ingestion of the data and processing of the data according to the processing schema.

(440) To address these issues, embodiments of the present disclosure relate to a data intake and query system that utilizes multiple components (e.g., connectors) to ingest (e.g., automatically ingest) data from one or more distributed ledger systems. The data intake and query system can initialize the components (e.g., the data intake and query system can include the components) and/or the data intake and query system can utilize components initialized by a separate system. While specific components may be referenced (e.g., connectors), it will be understood that any components may be utilized by the data intake and query system.

(441) All or a portion of the connectors can (e.g., automatically) obtain data from one or more distributed ledger systems. For example, the connectors can obtain data from nodes of a distributed ledger system. In some cases, the distributed ledger systems, the data intake and query system, and/or a separate system may include all or a portion of the nodes. The connector may, based on initialization of the particular connector, obtain all or a portion of the historical data associated with a distributed ledger system (e.g., data generated subsequent to the initialization of the distributed ledger system and prior to initialization of the particular connector) and may, periodically or aperiodically, obtain data generated subsequent to the connector previously obtaining data.

(442) To determine how to ingest the data from the distributed ledger systems, the data intake and query system can obtain processing schema(s) identifying one or more distributed ledger systems. For example, the data intake and query system can obtain the processing schema(s) during an initialization process. In response to obtaining the processing schema(s), the data intake and query system can initialize connectors to ingest data from the one or more distributed ledger systems identified by the processing schemas. The processing schema(s) may identify processing parameter(s), distributed ledger system(s), and/or computing device(s) for ingestion of the data. The data intake and query system can utilize the connectors to ingest data from the one or more distributed ledger systems using the processing schema(s). For all or a portion of the connectors,

the data intake and query system can assign a distributed ledger system (and/or a computing device) and a processing schema to the particular connector.

(443) In some cases, the processing schema(s) may include multiple processing schemas. For example, the processing schema(s) may include a first processing schema and a second processing schema. The processing schema(s) may indicate that the processing schema(s) are to be used to ingest data from the same or different distributed ledger systems. For example, the processing schema(s) may indicate that a first processing schema is to be used to ingest data from a first distributed ledger system and a second processing schema is to be used to ingest data from the first distributed ledger system or a second distributed ledger system.

(444) The data intake and query system can assign a respective processing schema of the multiple processing schemas to all or a portion of the connectors. For example, the data intake and query system can assign a first processing schema to a first connector and a second processing schema to a second connector. In some cases, one or more connectors may obtain the same data from the distributed ledger system (e.g., obtain the same data from the same nodes or different, parallel nodes associated with the distributed ledger system). Further, the connectors may process the data according to the particular processing schema assigned to the particular connector such that each connector may generate different processed data. For example, the first connector may process the data according to a first processing schema to generate first processed data and the second connector may process the data according to a second processing schema to generate second processed data.

(445) The data intake and query system can assign all or a portion of the connectors to ingest data from a distributed ledger system (e.g., different distributed ledger systems) for ingestion of data. For example, the data intake and query system can utilize one or more connectors for Ethereum and one or more connectors for Bitcoin.

(446) The data intake and query system can assign one or more connectors to ingest data from one or more computing devices (e.g., one or more non-distributed ledger systems). For example, the one or more computing devices may host, store data for, and/or be associated with one or more data sources (e.g., a threat database, a security database, a social media platform, etc.). In some cases, the data intake and query system can utilize one or more connectors (or other components) to ingest data associated with a URL, an IP address, and/or a machine identifier. Therefore, the data intake and query system can utilize connectors to ingest data from one or more distributed ledger systems and/or one or more computing devices.

(447) The connectors can process the data according to the processing schema assigned to the particular connector. For example, the connectors can enrich, filter, transform, tag, etc. the data according to the processing schema assigned to the particular connector.

(448) In some cases, the connectors and/or the data intake and query system can process data ingested from a first data source (e.g., a first distributed ledger system) based on data ingested from a second data source (e.g., a second distributed ledger system, a computing device, etc.). By processing data ingested from a first data source based on data ingested from a second data source, the connectors and/or the data intake and query system can combine, correlate, or otherwise utilize data obtained from multiple data sources (e.g., multiple distributed ledger system) to process the data.

(449) Based on ingesting the data from the distributed ledger system, the connectors and/or the data intake and query system may index the data. For example, the connectors and/or the data intake and query system may generate one or more virtual indices corresponding to the data. In response to receiving a query corresponding to the data from the distributed ledger system, the data intake and query system may utilize the one or more virtual indices corresponding to the data. In some cases, the connectors may index and store the data in a data store (e.g., a data base). For example, the connectors may index and store the data in a data store separate from the data intake and query system. The data intake and query system may utilize the stored data for query execution.

(450) The data intake and query system can cause display of the processed data obtained from the connectors (e.g., the processed data from the one or more distributed ledger systems and processed by the connectors). The processed data may include data associated with a single distributed ledger system, data associated with multiple distributed ledger systems, data associated with one or more distributed ledger systems and one or more computing devices, etc. The data intake and query system can cause display of the data via a GUI. For example, the data intake and query system may concurrently display data associated with a first distributed ledger in a first time-series visualization (e.g., a first chart), data associated with a second distributed ledger in a second time-series visualization (e.g., a second chart), and data associated with both the first distributed ledger and the second distributed ledger (e.g., a combination and/or correlation of the data associated with the first distributed ledger and the second distributed ledger) in a third time series visualization). By concurrently displaying data associated with multiple data sources (e.g., multiple distributed ledgers, a distributed ledger and a computing device, etc.), the data intake and query system can enable the comparison of the data from the different data sources, the tracing of data through one or more distributed ledger systems, the generation of predictive models based on distributed ledger data, the prediction of an output of a particular distributed ledger system (e.g., a predicted ledger transaction outcome), etc.

(451) Therefore, the use of a data intake and query system that utilizes multiple connectors to automatically ingest and index data from one or more distributed ledger systems and/or one or more computing devices enables the efficient execution of queries such that the data from the one or more distributed ledger systems and/or one or more computing devices need not be ingested in response to a query. For example, a device can generate a query for execution by the data intake and query system without reingesting the data the one or more distributed ledger systems and/or one or more computing devices.

(452) The data intake and query system can utilize connectors with different processing schemas to ingest and process data from the same distributed ledger system. By utilizing connectors with different processing schemas to ingest and process data from the same distributed ledger system, the data intake and query system can enable the efficient execution of queries. For example, a first connector may utilize a first processing schema to ingest and process data to generate first processed data that is responsive to a first subset of queries and a second connector may utilize a second processing schema to ingest and process data to generate second processed data that is responsive to a second subset of queries.

(453) The data intake and query system can concurrently display the data from multiple distributed ledger systems and/or multiple computing devices and correlations between the data. By concurrently displaying the data from multiple distributed ledger systems and/or multiple computing devices and correlations between the data, the data intake and query system can enable the concurrent display of such data without separate systems displaying the data via separate interfaces.

#### Operation of the Connectors for Multiple Systems

(454) Users may want to execute queries on data associated with distributed ledger systems. For example, as the use of distributed ledger systems continues to grow, a user of a data intake and query system may want to generate queries for execution on data from multiple distributed ledger systems (e.g., public distributed ledger systems). A user may want to provide the query and receive the query results associated with the data from the distributed ledger systems without separately ingesting data from the distributed ledger systems.

(455) The techniques described below can enable a data intake and query system to utilize connectors to ingest data from one or more distributed ledger systems and/or one or more computing devices, process the data, and display a result of the processed data. These techniques solve challenges of existing data intake and query system in that these systems may utilize connectors for multiple distributed ledger systems and/or connectors such that data can be

correlated and/or combined across distributed ledger systems.

(456) While existing system are able to execute queries on data from a distributed ledger system, such a query execution process can be inefficient for data associated with other queries and may not enable a user to correlate and/or combine data across distributed ledger systems. For example, the system may include one or more components that are customized and/or specific to a particular distributed ledger (e.g., Ethereum). As the one or more components may be customized and/or specific to a particular distributed ledger, the system may be unable to utilize the components for other distributed ledgers. Therefore, the system may be limited to executing queries on data associated with a particular distributed ledger system.

(457) The systems may be unable to correlate data associated with different distributed ledger systems and/or computing devices. For example, a client device may ingest data associated with a distributed ledger system and provide a query and these systems may execute the query on the ingested data to generate query results. Such query results may be inefficient and/or unresponsive as the query results may not identify correlations or combinations between data from different distributed ledger systems and/or computing devices and may require additional processing (e.g., an extensive and manual process). Therefore, the additional processing may be impractical for identification of correlations and/or combinations of data associated with different distributed ledger systems.

(458) The data intake and query system can utilize connectors to ingest (e.g., automatically ingest) data associated with multiple distributed ledger systems and/or computing systems. For example, the data intake and query system can utilize a first connector to ingest data associated with a first distributed ledger system and a second connector to ingest data associated with a second distributed ledger system. Each of the distributed ledger systems may generate data in different formats.

(459) To ingest the data from the distributed ledger systems, the data intake and query system may obtain processing schema(s). The processing schema(s) may identify one or more distributed ledger system(s) and/or computing device(s) and processing parameter(s) to obtain data from the distributed ledger system(s) and/or computing device(s). All or a portion of the distributed ledger system(s) and/or computing device(s) may be associated with a particular processing schema identifying how to ingest data from the particular distributed ledger system or computing device. Further, all or a portion of the distributed ledger system(s) and/or computing device(s) may be associated with multiple processing schemas for ingestion of data from the particular distributed ledger system or computing device.

(460) FIG. 27 is a block diagram of a data ingestion system **2700** that ingests data in accordance with example embodiments. As illustrated by FIG. 27, the data ingestion system **2700** includes a client device **204** in communication with a data intake and query system **108**, similar to that described above with reference to FIG. 18. The data intake and query system **108** receives an input from the client device **204** identifying a query for execution by the data intake and query system **108** on ingested data. The data ingestion system **2700** includes connectors **2702**, nodes **2704A** and **2704B**, distributed ledger systems **2706A** and **2706B**, and one or more computing devices **2708**. It will be understood that the data ingestion system **2700** may include more, less, or different distributed ledger systems and/or computing devices.

(461) In FIG. 27, the data ingestion system **2700** includes the distributed ledger systems **2706A** and **2706B** and the one or more computing devices **2708**. As discussed above, the distributed ledger systems **2706A** and **2706B** can provide one or more distributed ledgers or blockchains (e.g., public blockchains, private blockchains, consortium blockchains, hybrid blockchains, etc.) across one or more nodes **2704A**, **2704B** (e.g., full nodes, light nodes, archive nodes, etc.) that communicate via a network. For example, the distributed ledger systems **2706A** and **2706B** may include Ethereum, Bitcoin, Hyperledger Fabric, Quorum, Guardtime, KSI, Polkadot, Cosmos, Polygon, Arbitrum, etc. The distributed ledger systems **2706A** and **2706B** may be unrelated distributed ledger systems. For example, the distributed ledger system **2706A** may be Ethereum and the distributed ledger system

**2706B** may be Bitcoin.

(462) The distributed ledger systems **2706A** and **2706B** may each utilize one or more nodes that generate data. In the example of FIG. 27, the distributed ledger system **2706A** utilizes one or more nodes **2704A** and the distributed ledger system **2706B** utilizes one or more nodes **2704B**. In some cases, the one or more nodes are part of the data intake and query system **108**. In other cases, the one or more nodes are part of the distributed ledger systems **2706A** and **2706B** or a separate system.

(463) The one or more nodes **2704A** and the one or more nodes **2704B** may include one or more peer nodes, ordering nodes, etc. The peer nodes may be associated with a peer node blockchain (e.g., a blockchain that includes validated transactions and not non-validated transactions or blockchain that includes an indication of valid/invalid transactions). The ordering nodes may be associated with an ordering node blockchain (e.g., a blockchain that includes transactions regardless of validation).

(464) All or a portion of the one or more nodes **2704A** and the one or more nodes **2704B** can be implemented using individual computing devices, distributed processing systems, servers, isolated execution environments (e.g., containers, virtual machines, etc.), shared computing resources, etc. In some embodiments, the one or more nodes **2704A** and the one or more nodes **2704B** can be implemented on the same isolated execution environment or as part of different isolated execution environments (e.g., as different containers or pods of the same or different Kubernetes cluster or Docker swarm).

(465) The one or more nodes **2704A** and/or the one or more nodes **2704B** may be parallelized nodes. For example, the one or more nodes **2704A** may operate in parallel and perform the same data operations to generate the same data. Therefore, if one of the one or more nodes **2704A** is unavailable, the connectors **2702** can access the same data from a different one of the one or more nodes **2704A**.

(466) As discussed above, all or a portion of the one or more nodes **2704A** and the one or more nodes **2704B** may include one or more ledgers, which can be stored across one or more data stores. In some embodiments, the one or more ledgers **2704B** can include one or more blockchains.

(467) As discussed above, the one or more nodes **2704A** and the one or more nodes **2704B** can generate different types of data, including transactions, metrics data, log data, etc. For example, the one or more nodes **2704A**, **2704B** may generate price feed data. The different types of data generated by the one or more nodes **2704A** and the one or more nodes **2704B** and/or the distributed ledger systems **2706A** and **2706B** can be accessible via different paths or stored in different locations of the one or more nodes **2704A** and the one or more nodes **2704B**. For example, the data can be located in a data store, pub/sub, buffer, or real-time data stream. In some embodiments, such as when the distributed ledger system is implemented in an isolated execution environment system, the data can be wrapped or converted to another format, such as JSON, and stored as a JSON (or other type) file.

(468) The one or more computing devices **2708** may host, store, generate, or otherwise be associated with non-blockchain data (e.g., data that is not stored on a blockchain, offchain data, etc.). For example, the one or more computing devices **2708** may host, store, generate, or otherwise be associated with data that is associated with a blockchain but not stored on a blockchain (e.g., observed metric data associated with a blockchain, data identifying bad actors associated with a blockchain, etc.). The one or more computing devices **2708** may host, store, generate, or otherwise be associated with a threat database (e.g., a database identifying bad actors), a security database, a social media website, etc. For example, the one or more computing devices **2708** may host, store, generate, or otherwise be associated with Discord, Twitter, GitHub, etc.

(469) The data intake and query system **108** can ingest data from the distributed ledger systems **2706A** and **2706B** and the one or more computing devices **2708** for execution of one or more queries. The data intake and query system **108** can automatically ingest the data based on

processing schema(s) identifying the distributed ledger systems **2706A** and **2706B** and the one or more computing devices **2708**. Further, the data intake and query system **108** may periodically or aperiodically ingest updated data from the distributed ledger systems **2706A** and **2706B** and the one or more computing devices **2708**. In some cases, the data intake and query system **108** may ingest data from one or more servers, one or more databases, one or more resources, etc. For example, the data intake and query system **108** may utilize a resource identifier to ingest data from one or more servers.

(470) To ingest data from the distributed ledger systems **2706A** and **2706B** and the one or more computing devices **2708**, the data ingestion system **2700** may utilize connectors (e.g., distributed ledger connectors). For example, the data intake and query system **108** may include connectors **2702**. In some cases, the data intake and query system **108** may not include the connectors **2702** and the data intake and query system **108** may utilize connectors **2702** associated with another system. The data intake and query system **108** may initialize the connectors based on processing schema(s) to ingest data from the distributed ledger systems **2706A** and **2706B** and the one or more computing devices **2708**.

(471) All or a portion of the connectors **2702** may ingest data based on a processing schema. The processing schema may include instructions or processing parameters indicating how the connector is to process the data. For example, the processing schema may identify filters (e.g., filter criteria), enrichment parameters, tags, etc. In some cases, the processing schema may identify one or more operations to be performed by all or a portion of the connectors **2702** for ingestion of data (e.g., during ingestion of data).

(472) The data intake and query system **108** may assign a connector(s) **2702** to ingest data associated with a particular distributed ledger system **2706A** or **2706B** and/or computing devices **2708** based on the processing schema. In some embodiments, all or a portion of the connectors **2702** may be initialized to ingest data associated with a particular type of distributed ledger system. For example, a first connector may be initialized to ingest data associated with Ethereum and a second connector may be initialized to ingest data associated with Bitcoin.

(473) To ingest data from the distributed ledger systems **2706A** and **2706B**, the connectors **2702** may ingest data from the one or more nodes **2704A** of the distributed ledger system **2706A** and/or the one or more nodes **2704B** of the distributed ledger system **2706B**. The connectors **2702** can extract the data generated by the one or more nodes **2704A** and the one or more nodes **2704B** from a wrapper. Based on extracting the data generated by the one or more nodes **2704A** and the one or more nodes **2704B**, the connectors **2702** can store the data in one or more buckets of the data intake and query system **108**.

(474) The connectors **2702** can ingest data associated with the computing devices **2708**. For example, the connectors **2702** can ingest data generated by the computing devices **2708**. In another example, the connectors **2702** can ingest data from a data store, a cache, a database, etc. associated with the computing devices **2708**. In some cases, the data intake and query system **108** may utilize and/or may include separate components to ingest the data associated with the computing devices **2708**.

(475) As discussed above, the connectors **2702** may process the ingested data based on a processing schema assigned to the particular connector. For example, all or a portion of the connectors **2702** may transform, filter, enrich, or otherwise process the ingested data based on the processing schema assigned to the particular connector. In some cases, all or a portion of the connectors **2702** can enrich the ingested data with network-wide metrics, analytics (e.g., legend analytics), network data, and/or insights. For example, the connectors **2702** can enrich the ingested data with transaction cost, gas price (e.g., a cost to send a transaction for validation), block time, confirmed transactions, transactions to be discard, transactions to be validation, blockchain organization, etc.

(476) In some cases, all or a portion of the connectors **2702** may correlate, combine, compare, etc.

ingested data with data associated with one or more other distributed ledger systems (e.g., distributed ledger systems associated with different ledgers, different blockchain technologies, etc.) and/or computing devices. Specifically, a connector **2702** may correlate data ingested from computing devices **2708** with data ingested from one or more nodes **2704B** of distributed ledger system **2706B**. For example, a connector **2702** may correlate data (e.g., data identifying an address, account, wallet, etc. of an actor (a bad actor)) received from computing devices **2708** with data from the distributed ledger system **2706A** and may enrich data from the distributed ledger system **2706A** with the data received from the computing devices **2708**. Therefore, the connector **2702** can add a field, tag, etc. to the data from the distributed ledger system **2706A** identifying the data received from the computing devices **2708** (e.g., identifying a bad actor).

(477) In some cases, the data intake and query system **108** can process the ingested data. For example, the data intake and query system **108** can transform, filter, enrich, or otherwise process the ingested data. Further, the data intake and query system **108** may correlate, combine, compare, etc. data ingested from different distributed ledger systems (e.g., distributed ledger systems associated with different ledgers, different blockchain technologies, etc.) and/or computing devices.

(478) The data intake and query system **108** can store and index the processed, ingested data. For example, the data intake and query system **108** can store the processed, ingested data in one or more buckets and index the processed, ingested data for query execution.

(479) To index the data, the data intake and query system **108** can populate one or more inverted indices based on the ingested data. For example, the data intake and query system **108** can populate a first inverted index based on data received from the distributed ledger system **2706A** and a second inverted index based on data received from the distributed ledger system **2706B**. As discussed above, the data intake and query system **108** may utilize the connectors **2702** to extract data (e.g., a transaction identifier, a node identifier, channel identifier, ledger identifier, ledger start time, etc.) from the data from the distributed ledger systems **2706A** and **2706B**. The data intake and query system **108** can include the extracted data as keywords or field-value pairs and store the keywords or field-value pairs in one or more inverted indices, such as the inverted index described herein with reference to FIG. 5B.

(480) In some cases, the inverted indices may be “virtual indices.” Therefore, the data intake and query system **108** can implement a virtual index based on the data ingested from the particular distributed ledger system **2706A** or **2706B** and/or computing devices **2708**. For example, the data intake and query system **108** can implement a first virtual index for the distributed ledger system **2706A**, a second virtual index for the distributed ledger system **2706B**, a third virtual index for the computing devices **2708**, etc.

(481) The data intake and query system **108** can obtain queries (e.g., from the client device **204**) associated with the ingested, processed data. As the connectors **2702** and/or the data intake and query system **108** can compare, correlate, combine, bridge, contrast, etc. data associated with the distributed ledger systems **2706A** and/or **2706B** and/or computing devices **2708**, the data intake and query system **108** can execute a single instance of the query across data associated with and/or correlated across multiple distributed ledger systems and/or computing devices. For example, the distributed ledger system **2706A** and the distributed ledger system **2706B** may interoperate and a transaction on the distributed ledger system **2706A** may be bridged with a transaction on the distributed ledger system **2706B**. Therefore, the data intake and query system **108** and/or the connectors may correlate the transactions for efficient execution of queries.

(482) The data intake and query system **108** may implement an External Result Provider (“ERP”) process to provide an interface through which a search head may access the virtual indices for execution of a query. The ERP process can connect to respective virtual indices. The search head may spawn only a single ERP process in response to multiple virtual indexes referenced in a search request, or the search head may spawn different ERP processes for different virtual indices.

(483) The ERP process may generate search requests in the protocol and syntax of the respective



virtual indices. Upon receiving search results from their corresponding virtual indices, the respective ERP process can pass the result to the search head, which may return the results to a core engine for processing and may terminate operation of the respective ERP process.

(484) The data intake and query system **108** can cause display of a visualization of the ingested data (e.g., the ingested, processed data) via an interface (e.g., a GUI) of the client device **204**. In some cases, the data intake and query system **108** can cause display of the visualization via an interface of a different computing device. The visualization may include a visualization of the ingested data (e.g., a visualization of one or more transactions on a blockchain), a visualization comparing, contrasting, correlating, combining, etc. ingested data associated with multiple distributed ledger systems and/or computing devices.

#### Operation of Versioned Connectors for a System

(485) As discussed above, users may want to execute multiple queries on data associated with the same distributed ledger system. A user may want to provide the query and receive the query results associated with the data from the distributed ledger system without separately ingesting data from the distributed ledger system. It may be advantageous to automatically ingest data a distributed ledger system using multiple processing schemas.

(486) The techniques described below can enable a data intake and query system to utilize multiple connectors to ingest data from nodes of a distributed ledger system using multiple processing schemas. These techniques solve challenges of existing data intake and query system in that these systems may utilize connectors and associated processing schemas to ingest the same data. While existing system are able to execute queries on data from a distributed ledger system, such a query execution process can be inefficient for multiple queries and/or diverse queries. For example, existing systems may require a first user to ingest data from a distributed ledger system for querying and a second user to ingest the data from the distributed ledger system. Further, the system may ingest data in response to a received search query according to the parameters of a search query. Therefore, the system may be limited to ingesting data associated with a distributed ledger system as part of the query execution.

(487) To ingest the data from the distributed ledger system, the data intake and query system can obtain processing schemas. The data intake and query system can initialize multiple connectors to ingest data from a distributed ledger system. All or a portion of the multiple connectors may ingest the data according to a particular processing schema or a version of a shared processing schema. Therefore, the multiple “versioned” connectors may ingest data (e.g., the same data) and generate ingested, processed data (e.g., different ingested, processed data).

(488) FIG. **28** is a block diagram of a data ingestion system **2800** that ingests data in accordance with example embodiments. As illustrated by FIG. **28**, the data ingestion system **2800** includes a client device **204** in communication with a data intake and query system **108**, similar to that described above with reference to FIG. **18** and FIG. **27**. The data ingestion system **2800** includes connectors **2802A**, **2802B**, . . . , **2802N**, nodes **2806**, and distributed ledger system **2706A**. It will be understood that the data ingestion system **2800** may include more, less, or different distributed ledger systems and/or computing devices.

(489) In FIG. **28**, the data ingestion system **2800** includes the distributed ledger system **2706A**. As discussed above, the distributed ledger system **2706A** can provide one or more distributed ledgers or blockchains across one or more nodes **2806** that communicate via a network.

(490) As discussed above, the data intake and query system **108** can ingest data from the distributed ledger system **2706A** for execution of one or more queries. The data intake and query system **108** can automatically ingest the data based on the processing schemas **2804A**, **2804B**, . . . , **2804N**. For example, the processing schemas **2804A**, **2804B**, . . . , **2804N** may indicate data sources for ingestion of data into the data intake and query system **108**. Further, the processing schemas **2804A**, **2804B**, . . . , **2804N** may indicate processing parameters (e.g., filters, transformation parameters, enrichment parameters, tags, etc.) for processing of the ingested data. In some cases, all

or a portion of the processing schemas **2804A**, **2804B**, . . . , **2804N** may correspond to different versions of a shared set of processing parameters.

(491) Based on the processing schemas **2804A**, **2804B**, . . . , **2804N**, the data intake and query system **108** may periodically or aperiodically ingest updated data from the distributed ledger system **2706A**. For example, the data intake and query system **108** may ingest data identifying historical transactions associated with the distributed ledger system **2706A** and may ingest data identifying subsequent transactions associated with the distributed ledger system **2706A**.

(492) As discussed above, to ingest data from the distributed ledger system **2706A**, the data ingestion system **2700** may utilize connectors **2802A**, **2802B**, . . . , **2802N** (e.g., distributed ledger connectors). In the example of FIG. **28**, the data intake and query system **108** does not include the connectors **2802A**, **2802B**, . . . , **2802N**. However, it will be understood that the data intake and query system **108** may or may not include the connectors **2802A**, **2802B**, . . . , **2802N**. The data intake and query system **108** may include more, less, or different connectors.

(493) The data intake and query system **108** may assign a respective processing schema **2804A**, **2804B**, . . . , **2804N** to all or a portion of the connectors **2802A**, **2802B**, . . . , **2802N**. The processing schemas **2804A**, **2804B**, . . . , **2804N** may indicate the nodes **2806**, the distributed ledger system **2706A**, and/or processing parameters for ingestion of data. In the example of FIG. **28**, the data intake and query system **108** assigns the processing schema **2804A** to the connector **2802A**, the processing schema **2804B** to the connector **2802B**, and the processing schema **2804N** to the connector **2802N**.

(494) The processing schemas **2804A**, **2804B**, . . . , **2804N** may identify the same nodes **2806** and/or the same distributed ledger system **2706A**. For example, the processing schemas **2804A**, **2804B**, . . . , **2804N** may identify different parallelized nodes of the same distributed ledger system **2706A** that each generate the same or similar data for ingestion (e.g., the parallelized nodes may generate data with different hosts). However, the processing schemas **2804A**, **2804B**, . . . , **2804N** may identify different processing parameters for processing the data. For example, all or a portion of the processing schemas **2804A**, **2804B**, . . . , **2804N** may identify different filters, transformation parameters, enrichment parameters, tags, etc. to be applied to the data (e.g., during data ingestion). Therefore, all or a portion of the connectors **2802A**, **2802B**, . . . , **2802N** may ingest data from the nodes **2806** (e.g., the same data) and apply different processing parameters to the ingested data to generate different ingested, processed data.

(495) The assignment of processing schemas **2804A**, **2804B**, . . . , **2804N** to connectors **2802A**, **2802B**, . . . , **2802N** and/or connectors **2802A**, **2802B**, . . . , **2802N** to nodes **2806** may be a dynamic assignment. For example, the data intake and query system **108** may dynamically assign (or map) a processing schema to a particular connector and/or a particular connector to a particular node. The data intake and query system **108** may obtain updated processing schemas, new processing schemas, discontinued processing schemas, etc. and may assign the processing schemas to the connectors (e.g., by updating a prior assignment of the processing schemas to the connectors).

(496) In some cases, the assignment of processing schemas **2804A**, **2804B**, . . . , **2804N** to connectors **2802A**, **2802B**, . . . , **2802N** and/or connectors **2802A**, **2802B**, . . . , **2802N** to nodes **2806** may be based on observed or monitored query characteristics. For example, the data intake and query system **108** (or a separate system) may monitor received queries. The data intake and query system **108** may identify particular queries, particular query parameters, etc. based on monitoring the received queries. Further, the data intake and query system **108** may identify particular queries, particular query parameters, etc. associated with a particular distributed ledger system, a group of distributed ledger systems, etc.

(497) The data intake and query system **108** may identify particular queries, particular query parameters, etc. that are associated with a threshold percentage or a threshold number of received queries. For example, the data intake and query system **108** may determine that a particular query

or query parameter (e.g., a query associated with a number of transactions, a gas price, a number of transactions in a particular time period, a type of transactions, articles related to a particular subject or headline, bad actor data, etc.) is associated with 50% of queries associated with a particular distributed ledger system. In response to determining that the number of queries associated with a particular query or query parameter exceeds or matches a threshold value (e.g., 50%, 75%, etc.), the data intake and query system **108** can provide a processing schema associated with the particular query or query parameter to a connector. Therefore, the connector can automatically ingest data according to the particular query or query parameter. By automatically ingesting data according to a particular query or query parameter, the responsiveness and/or efficiency of the data intake and query system can be greatly increased as compared to traditional data intake and query systems. In some cases, for a user (e.g., a new user), the data intake and query system **108** may dynamically integrate the automatic ingestion to automatically ingest data according to a particular query (ies) (e.g., a set of queries associated with more users than one or more other queries).

(498) In another example, the data intake and query system **108** may determine that 50% of queries are associated with a particular distributed ledger system. In response to determining that the number of queries associated with a distributed ledger system exceeds or matches a threshold value (e.g., 50%, 75%, etc.), the data intake and query system **108** can assign a connector to the particular distributed ledger system.

(499) The data intake and query system **108** can manage the assignment of processing schemas **2804A, 2804B, . . . , 2804N** to connectors **2802A, 2802B, . . . , 2802N**. For example, the data intake and query system **108** can perform an initial assignment of processing schemas **2804A, 2804B, . . . , 2804N** to connectors **2802A, 2802B, . . . , 2802N**. To manage the assignment of the processing schemas **2804A, 2804B, . . . , 2804N** to connectors **2802A, 2802B, . . . , 2802N**, subsequent to the initial assignment, the data intake and query system **108** can perform one or more operations. For example, the one or more operations can include assignment of a different processing schema to a particular connector, updating of a processing schema assigned to a particular connector, discontinuing use of a particular connector associated with a particular processing schema, initialization of a connector and assignment of a processing schema to the initialized connector, etc.

(500) The data intake and query system **108** can manage the assignment of processing schemas **2804A, 2804B, . . . , 2804N** to connectors **2802A, 2802B, . . . , 2802N** based on and/or in response to observed or monitored query characteristics. For example, the data intake and query system **108** (or a separate system) may monitor received queries. The data intake and query system **108** may identify particular queries, particular query parameters, etc. based on monitoring the received queries. The data intake and query system **108** may perform the one or more operations to manage the assignment based on monitoring the received queries.

(501) Based on the assignment of processing schemas **2804A, 2804B, . . . , 2804N** to connectors **2802A, 2802B, . . . , 2802N**, all or a portion of the connectors **2802A, 2802B, . . . , 2802N** may ingest data from the nodes **2806**. Therefore, all or a portion of the connectors **2802A, 2802B, . . . , 2802N** may ingest data associated with the distributed ledger system **2706A**. As discussed above, to ingest the data, the connectors **2802A, 2802B, . . . , 2802N** can extract the data generated by the nodes **2806**. The connectors **2802A, 2802B, . . . , 2802N** may ingest and process the data generated by the nodes **2806** based on the processing schemas **2804A, 2804B, . . . , 2804N**. Therefore, the connectors **2802A, 2802B, . . . , 2802N** may generate ingested, processed data.

(502) All or a portion of the connectors **2802A, 2802B, . . . , 2802N** may generate respective ingested, processed data based on the processing schemas **2804A, 2804B, . . . , 2804N**. All or a portion of the connectors **2802A, 2802B, . . . , 2802N** can ingest data and process the data by filtering the data, tagging the data (e.g., tagging a transaction as bridged with another transaction associated with another distributed ledger system), modifying and/or verifying a data format of the data, transforming the data, enriching the data, etc. according to the processing schema assigned to

the particular connector. For example, a first connector may obtain data from one or more first parallel nodes of the nodes **2806** and a second connector may obtain the same or different data from one or more second parallel nodes of the nodes **2806**. Further, the first connector may provide ingested, processed data with a first data format and/or including first data and the second connector may provide ingested, processed data with a second data format and/or including second data. Based on extracting the data generated by the nodes **2806**, the connectors **2802A**, **2802B**, . . . , **2802N** can store the data in one or more buckets of the data intake and query system **108** for the execution of queries.

(503) In some cases, as all or a portion of the connectors **2802A**, **2802B**, . . . , **2802N** may be ingesting the same data from the nodes **2806** of the distributed ledger system **2706A**, one or more portions of the data may be duplicated during ingestion. For example, all or a portion of the ingested, processed data from the connector **2802A**, the ingested, processed data from the connector **2802B**, and/or the ingested, processed data from the connector **2802N** may be duplicated. The data intake and query system **108** may perform one or more deduplication operations to remove the duplicated data from the one or more buckets (e.g., remove data such that a single copy of the data is stored in the one or more buckets).

(504) In some cases, a connector may ingest data prior to other collectors. For example, a connector may be assigned to ingest data from a node that has a greater number of peer nodes than other nodes of the nodes **2806**. Based on ingesting the data prior to other collectors, the data ingested by the connector may be retained as the single copy. In some cases, the data intake and query system **108** may not remove duplicated data.

(505) In some embodiments, the data intake and query system **108** may include (implement, initialize, etc.) the nodes of the distributed ledger system **2706A**. With reference to FIG. **29**, FIG. **29** is a block diagram of a data ingestion system **2900** that ingests data in accordance with example embodiments. As illustrated by FIG. **29**, the data ingestion system **2800** includes a client device **204** in communication with a data intake and query system **108**, similar to that described above with reference to FIG. **18**. The data intake and query system **108** includes connectors **2802A**, **2802B**, . . . , **2802N**, and nodes **2806**. The nodes **2806** may be in communication with the distributed ledger system **2706A** to generate data. It will be understood that the data ingestion system **2900** may include more, less, or different distributed ledger systems and/or computing devices.

(506) In FIG. **29**, the data intake and query system **108** includes the connectors **2802A**, **2802B**, . . . , **2802N**, and the nodes **2806**. Therefore, the data intake and query system **108** can automatically ingest data from distributed ledger system **2706A** using nodes **2806** that are internally managed by the data intake and query system **108**.

#### User Interfaces

(507) Users may want to generate customized dashboards based on data ingested from one or more distributed ledger systems and/or one or more computing devices by a data intake and query system. For example, a user may want to customize a dashboard to illustrate changes in a price feed. It may be advantageous to modify how a user can customize the dashboard for display. The techniques described below can enable a user to provide a custom dashboard based on ingested data. These techniques solve the challenges of existing systems in that these techniques may enable users to provide comparisons, correlations, etc. between data across multiple distributed ledger systems and/or computing devices. While the existing systems may enable users to provide a visualization of data associated with a single distributed ledger system, the systems may not enable users to provide a visualization of data associated with multiple distributed ledger systems and/or computing devices. To obtain data associated with multiple distributed ledger systems and/or computing devices, enrich the data, and provide a visualization of the data, the system generally requires an extensive and inefficient manual processes and often the user is unable to customize a dashboard for display of the data.

(508) In the presently disclosed system, a data intake and query system (or a separate system) can ingest data associated with multiple distributed ledger systems and/or computing devices using one or more connectors. Further, the data intake and query system can define parameters for the generation of customized visualizations. This data intake and query system enable an efficient and consistent customized dashboard generation process. The data intake and query system may cause the implementation of a GUI to provide the selectable parameters for selection to the client device. For example, the data intake and query system may cause a GUI to be displayed on the client device and may populate the GUI with the selectable parameters. The selectable parameters may include different criteria, parameters, filters, etc. that identify the data for display.

(509) FIG. 30 illustrates an example interface 3000 showing various exemplary features in accordance with one or more embodiments. In the illustrated embodiment of FIG. 30, the interface 3000 includes a definition section 3010, a data source selection section 3002, a data source parameter section 3004, a parameter comparison section 3008, a parameter combination section 3009, a comparison results section 3012, a combination results section 3013, a first results section 3020, and a second results section 3030. The interface 3000 may include various implementation elements including a first implementation element 3014 and a second implementation element 3016. It will be understood that the interface 3000 may include, more, less, or different elements.

(510) The interface 3000 may be an exemplary configuration screen showing selectable parameters for the generation of a customized dashboard. In some embodiments, the interface 3000 can include a hide options interface element that enables a user to toggle between hiding certain options or selectors and showing the options or selectors. In this way, the system can use less real estate on a screen. The example interface 3000 is illustrative of an interface that a computing system (e.g., a server in communication with the metric analyzer) generates and presents to a user. In the example of FIG. 30, the interface 3000 includes a particular selection of selectable parameters to customize a dashboard (e.g., for display to a user). For example, the generation of the customized dashboard may be based on an interaction by the user with one or more of the data source selection section 3002, the data source parameter section 3004, the parameter comparison section 3008, or the parameter combination section 3009. In some embodiments, all or a portion of the data source selection section 3002, the data source parameter section 3004, the parameter comparison section 3008, or the parameter combination section 3009 may be populated with selectable parameters by the data intake and query system.

(511) As a non-limiting example, one or more elements of the definition section 3010 can be used to identify data for the generation of a customized dashboard. In response to a request for generation of dashboard, the data intake and query system may generate the dashboard based on a selection of the selectable parameters. The generated dashboard can be displayed via an interactive display and may include the comparison results section 3012, the combination results section 3013, the first results section 3020, and/or the second results section 3030.

(512) The data source selection section 3002 can accept user input in the form of a selection of a particular data source based on one or more search terms, keywords, or tokens. The keywords can include a search string, text, selections from drop down menus, and/or interaction with or movement of interactive user interface elements, or other inputs that can be used to identify a data source. In response to the input received via the data source selection section 3002, the data intake and query system can identify data sources (e.g., one or more distributed ledger systems, computing devices, etc.) and associated ingested data for inclusion in the dashboard.

(513) The data source parameter section 3004 can accept user input in the form of a selection of a particular parameter (e.g., a value, a key, a key-value pair, etc.) of the data ingested from one or more data sources based on one or more search terms, keywords, or tokens. The keywords can include a search string, text, selections from drop down menus, and/or interaction with or movement of interactive user interface elements, or other inputs that can be used to identify a data source parameter. In response to the input received via the data source parameter section 3004, the

data intake and query system can identify parameters for inclusion in the dashboard.

(514) The parameter comparison section **3008** can accept user input in the form of a selection of one or more parameters of the selected data sources for comparison based on one or more search terms, keywords, or tokens. The keywords can include a search string, text, selections from drop down menus, and/or interaction with or movement of interactive user interface elements, or other inputs. The selected parameters can identify how the data intake and query system is to compare data from multiple distributed ledger systems and/or computing devices. The data intake and query system may populate the comparison results section **3012** of the dashboard with a preview of the results of comparing the data from the multiple distributed ledger systems and/or computing devices.

(515) The parameter combination section **3009** can accept user input in the form of a selection of one or more parameters of the selected data sources for combination based on one or more search terms, keywords, or tokens. The keywords can include a search string, text, selections from drop down menus, and/or interaction with or movement of interactive user interface elements, or other inputs. The selected parameters can identify how the data intake and query system is to combine data from multiple distributed ledger systems and/or computing devices. The data intake and query system may populate the combination results section **3013** of the dashboard with a preview of the results of combining the data from the multiple distributed ledger systems and/or computing devices.

(516) In some cases, one or more of the data source selection section **3002**, the data source parameter section **3004**, the parameter comparison section **3008**, or the parameter combination section **3009** can be left blank. For example, the parameter combination section **3009** can be left blank such that the dashboard does not include a combination results section **3013**. Similarly, other fields or selectors within the definition section **3010** may be left in a blank or null state.

Accordingly, in certain cases, if a parameter is left empty, blank, or in a null state, the system can return results related to the empty selector.

(517) In the example of FIG. 30, data source selection section **3002** identifies the selected data sources as “All” (e.g., all distributed ledger systems), the data source parameter section **3004** identifies the selected parameters as “All,” the parameter comparison section **3008** identifies the parameters for comparison as “Activities” (e.g., the number of activities or transactions), and the parameter combination section **3009** identifies the parameters for combination as “Activities.” Further, the comparison results section **3012** identifies that “Blockchain #1” is associated with “56” activities, “Blockchain #10” is associated with “41” activities, “Non-Blockchain #1” is associated with “19” activities, “Non-Blockchain #2” is associated with “8” activities, and “Blockchain #8” is associated with “5” activities. The combination results section **3013** identifies that the selected data sources are associated with “319” total activities. It will be understood that the definition section **3010** may include any parameters and/or data sources and the user may specify any value for the parameters and/or data sources.

(518) The interface **3000** may include a first implementation element **3014**. The first implementation element **3014** can enable a user to apply the parameters defined in the definition section **3010** for the generation of a customized dashboard. Further, the interface **3000** may include a second implementation element **3016**. The second implementation element **3016** can enable a user to reset the parameters defined in the definition section **3010**.

(519) The first results section **3020** may include a preview of one or more graphs or representations of the data ingested from one or more distributed ledger systems and/or computing devices. For example, the first results section **3020** may include a bar graph, a line graph, a symbolical representation, a numerical summarization, a textual representation, or any representation or summarization of the data. In the example of FIG. 30, the first results section **3020** includes a bar graph.

(520) A second results section **3030** may include a preview of data ingested from one or more

distributed ledger systems and/or computing devices for display via the dashboard. The second results section **3030** may identify a time **3024** (e.g., a time of an activity or a transaction) and source **3026** (e.g., a source of the activity or transaction). It will be understood that the second results section **3030** can be associated with more, less, or different information as desired. In some cases, the second results section **3030** may update in real time as data is ingested.

(521) In some cases, the interface **3000** may include a query element to enable users to generate queries for execution on data associated with multiple distributed ledger systems and/or computing devices. For example, the interface **3000** may include a query element that enables users to define queries for execution. In some embodiments, the data intake and query system (or a separate system) may obtain one or more defined queries associated with a particular user. The data intake and query system may process the one or more defined queries to generate a query recommendation for the same user or a different user. For example, a first user and a second user may be associated with the same distributed ledger and the data intake and query system may, for a second user, recommend (e.g., by causing display of) a query defined by the first user for performance.

#### Versioned Connectors for Data Ingestion

(522) As discussed above, data intake and query system can utilize connectors to (e.g., automatically) obtain data from one or more distributed ledger systems and/or one or more computing devices. The data intake and query system can identify a processing schema(s) for ingestion of the data from one or more distributed ledger systems. The data intake and query system can utilize the connectors to ingest data from the one or more distributed ledger systems using the processing schema(s). Therefore, for all or a portion of the connectors, the data intake and query system can assign a distributed ledger system (and/or a computing device) and a processing schema to the particular connector. The connectors may process the data according to the particular processing schema assigned to the particular connector such that each connector may generate different processed data. In some cases, the connectors and/or the data intake and query system can process data ingested from a first data source (e.g., a first distributed ledger system) based on data ingested from a second data source (e.g., a second distributed ledger system, a computing device, etc.).

(523) With reference to FIG. **31**, an illustrative algorithm or routine **3100** will be described for implementing connectors to ingest data from one or more distributed ledger systems and/or computing devices. The routine **3100** may be implemented, for example, by the data intake and query system **108** described above with reference to FIG. **2**. The routine **3100** begins at block **3102**, where the data intake and query system **108** receives blockchain data. For example, the blockchain data may be data from one or more distributed ledger systems (e.g., the blockchain data may be associated with a blockchain network). The blockchain data may be generated by a plurality of nodes of the one or more distributed ledger systems and/or of the data intake and query system **108** (e.g., nodes internal and/or external to the data intake and query system **108**). For example, at least a portion of the blockchain data may be based on transactions between the plurality of nodes. At least a portion of the blockchain data may include transaction data for a blockchain transaction between multiple nodes. In some cases, the data intake and query system **108** can receive non-blockchain data generated by a computing device (e.g., data not stored as part of a blockchain, offchain data, etc.). Therefore, the data intake and query system **108** can receive the blockchain data.

(524) The data intake and query system **108** may identify (e.g., generate, obtain, etc.) an assignment of processing schemas to connectors. The data intake and query system **108** may manage the assignment by updating a processing schema assigned to a particular connector for processing of additional blockchain data, removing or discontinuing use of a particular processing schema, and/or adding a processing schema for processing of additional blockchain data.

(525) To ingest the data, at block **3104**, the data intake and query system **108** processes the

blockchain data according to a first processing schema. Processing the blockchain data according to the first processing schema may include processing the blockchain data using a first set of filters or enrichment parameters associated with the first processing schema. The data intake and query system **108** may process the blockchain data using a first connector that is assigned the first processing schema. In some embodiments, the data intake and query system **108** may process additional blockchain data (e.g., associated with a different blockchain network and/or generated by an additional plurality of nodes) using the first connector and the first processing schema. In other embodiments, the data intake and query system **108** may process the additional blockchain data using the first connector and a third processing schema assigned to the additional blockchain data. Therefore, the data intake and query system **108** can process the blockchain data according to the first processing schema.

(526) At block **3106**, the data intake and query system **108** processes the blockchain data according to a second processing schema. Processing the blockchain data according to the second processing schema may include processing the blockchain data using a second set of filters or enrichment parameters associated with the second processing schema. The data intake and query system **108** may process the blockchain data using a second connector that is assigned the second processing schema. In some embodiments, the data intake and query system **108** may process additional blockchain data using the second connector and the second processing schema. In other embodiments, the data intake and query system **108** may process the additional blockchain data using the second connector and a fourth processing schema assigned to the additional blockchain data. The first connector and the second connector may process the blockchain data based on characteristics of the blockchain network associated with the blockchain data. In some cases, the data intake and query system **108** may process non-blockchain data to generate processed non-blockchain data. Therefore, the data intake and query system **108** can process the blockchain data according to the second processing schema.

(527) In some cases, the data intake and query system **108** can perform operations on data associated with multiple blockchain networks (e.g., the data and the additional data) and/or data from non-blockchain data sources. For example, the data intake and query system can correlate, compare, combine, contrast, etc. first blockchain data associated with a first blockchain network and at least one of second blockchain data associated with a second blockchain network or non-blockchain data generated by a computing device.

(528) Based on processing the blockchain data, at block **3108**, the data intake and query system **108** routes the processed blockchain data. For example, the data intake and query system **108** may route the processed blockchain data, via the first connector and the second connector, to a component of the data intake and query system **108** (e.g., one or more buckets). The processed blockchain data may include first processed blockchain data (processed using the first processing schema) and second processed blockchain data (processed using the second processing schema). The data intake and query system **108** may route multiple sets of processed blockchain data associated with multiple distributed ledger systems (e.g., multiple blockchain networks) and/or processed non-blockchain data. In some cases, the data intake and query system **108** may generate one or more virtual indices associated with the processed blockchain data (e.g., a first virtual index associated with the first processed blockchain data and a second virtual index associated with the second processed blockchain data). Therefore, the data intake and query system **108** can route the processed blockchain data.

(529) In some cases, the data intake and query system **108** can filter the processed blockchain data to remove duplicated data. For example, the data intake and query system **108** may perform a deduplication operation.

(530) The data intake and query system **108** may cause display, via a display device, of an interface that includes visualizations associated with the blockchain network (e.g., of the processed blockchain data). The visualizations may be generated based on the processed blockchain data. In



some cases, the visualizations may indicate a blockchain status and/or operational device status of at least one node.

## Terminology

(531) Computer programs typically comprise one or more instructions set at various times in various memory devices of a computing device, which, when read and executed by at least one processor, will cause a computing device to execute functions involving the disclosed techniques. In some embodiments, a carrier containing the aforementioned computer program product is provided. The carrier is one of an electronic signal, an optical signal, a radio signal, or a non-transitory computer-readable storage medium.

(532) Any or all of the features and functions described above can be combined with each other, except to the extent it may be otherwise stated above or to the extent that any such embodiments may be incompatible by virtue of their function or structure, as will be apparent to persons of ordinary skill in the art. Unless contrary to physical possibility, it is envisioned that (i) the methods/steps described herein may be performed in any sequence and/or in any combination, and (ii) the components of respective embodiments may be combined in any manner.

(533) Although the subject matter has been described in language specific to structural features and/or acts, it is to be understood that the subject matter defined in the appended claims is not necessarily limited to the specific features or acts described above. Rather, the specific features and acts described above are disclosed as examples of implementing the claims, and other equivalent features and acts are intended to be within the scope of the claims.

(534) Conditional language, such as, among others, “can,” “could,” “might,” or “may,” unless specifically stated otherwise, or otherwise understood within the context as used, is generally intended to convey that certain embodiments include, while other embodiments do not include, certain features, elements and/or steps. Thus, such conditional language is not generally intended to imply that features, elements and/or steps are in any way required for one or more embodiments or that one or more embodiments necessarily include logic for deciding, with or without user input or prompting, whether these features, elements and/or steps are included or are to be performed in any particular embodiment.

(535) Unless the context clearly requires otherwise, throughout the description and the claims, the words “comprise,” “comprising,” and the like are to be construed in an inclusive sense, as opposed to an exclusive or exhaustive sense, i.e., in the sense of “including, but not limited to.” As used herein, the terms “connected,” “coupled,” or any variant thereof means any connection or coupling, either direct or indirect, between two or more elements; the coupling or connection between the elements can be physical, logical, or a combination thereof. Additionally, the words “herein,” “above,” “below,” and words of similar import, when used in this application, refer to this application as a whole and not to any particular portions of this application. Where the context permits, words using the singular or plural number may also include the plural or singular number respectively. The word “or” in reference to a list of two or more items, covers all of the following interpretations of the word: any one of the items in the list, all of the items in the list, and any combination of the items in the list. Likewise, the term “and/or” in reference to a list of two or more items, covers all of the following interpretations of the word: any one of the items in the list, all of the items in the list, and any combination of the items in the list.

(536) Conjunctive language such as the phrase “at least one of X, Y and Z,” unless specifically stated otherwise, is otherwise understood with the context as used in general to convey that an item, term, etc. may be either X, Y or Z, or any combination thereof. Thus, such conjunctive language is not generally intended to imply that certain embodiments require at least one of X, at least one of Y and at least one of Z to each be present. Further, use of the phrase “at least one of X, Y or Z” as used in general is to convey that an item, term, etc. may be either X, Y or Z, or any combination thereof.

(537) In some embodiments, certain operations, acts, events, or functions of any of the algorithms

described herein can be performed in a different sequence, can be added, merged, or left out altogether (e.g., not all are necessary for the practice of the algorithms). In certain embodiments, operations, acts, functions, or events can be performed concurrently, e.g., through multi-threaded processing, interrupt processing, or multiple processors or processor cores or on other parallel architectures, rather than sequentially.

(538) Systems and modules described herein may comprise software, firmware, hardware, or any combination(s) of software, firmware, or hardware suitable for the purposes described. Software and other modules may reside and execute on servers, workstations, personal computers, computerized tablets, PDAs, and other computing devices suitable for the purposes described herein. Software and other modules may be accessible via local computer memory, via a network, via a browser, or via other means suitable for the purposes described herein. Data structures described herein may comprise computer files, variables, programming arrays, programming structures, or any electronic information storage schemes or methods, or any combinations thereof, suitable for the purposes described herein. User interface elements described herein may comprise elements from graphical user interfaces, interactive voice response, command line interfaces, and other suitable interfaces.

(539) Further, processing of the various components of the illustrated systems can be distributed across multiple machines, networks, and other computing resources. In certain embodiments, one or more of the components of the data intake and query system **108** can be implemented in a remote distributed computing system. In this context, a remote distributed computing system or cloud-based service can refer to a service hosted by one more computing resources that are accessible to end users over a network, for example, by using a web browser or other application on a client device to interface with the remote computing resources. For example, a service provider may provide a data intake and query system **108** by managing computing resources configured to implement various aspects of the system (e.g., search head **210**, indexers **206**, etc.) and by providing access to the system to end users via a network.

(540) When implemented as a cloud-based service, various components of the system **108** can be implemented using containerization or operating-system-level virtualization, or other virtualization technique. For example, one or more components of the system **108** (e.g., search head **210**, indexers **206**, etc.) can be implemented as separate software containers or container instances. Each container instance can have certain resources (e.g., memory, processor, etc.) of the underlying host computing system assigned to it, but may share the same operating system and may use the operating system's system call interface. Each container may provide an isolated execution environment on the host system, such as by providing a memory space of the host system that is logically isolated from memory space of other containers. Further, each container may run the same or different computer applications concurrently or separately, and may interact with each other. Although reference is made herein to containerization and container instances, it will be understood that other virtualization techniques can be used. For example, the components can be implemented using virtual machines using full virtualization or paravirtualization, etc. Thus, where reference is made to “containerized” components, it should be understood that such components may additionally or alternatively be implemented in other isolated execution environments, such as a virtual machine environment.

(541) Likewise, the data repositories shown can represent physical and/or logical data storage, including, e.g., storage area networks or other distributed storage systems. Moreover, in some embodiments the connections between the components shown represent possible paths of data flow, rather than actual connections between hardware. While some examples of possible connections are shown, any of the subset of the components shown can communicate with any other subset of components in various implementations.

(542) Embodiments are also described above with reference to flow chart illustrations and/or block diagrams of methods, apparatus (systems) and computer program products. Each block of the flow

chart illustrations and/or block diagrams, and combinations of blocks in the flow chart illustrations and/or block diagrams, may be implemented by computer program instructions. Such instructions may be provided to a processor of a general purpose computer, special purpose computer, specially-equipped computer (e.g., comprising a high-performance database server, a graphics subsystem, etc.) or other programmable data processing apparatus to produce a machine, such that the instructions, which execute via the processor(s) of the computer or other programmable data processing apparatus, create means for implementing the acts specified in the flow chart and/or block diagram block or blocks. These computer program instructions may also be stored in a non-transitory computer-readable memory that can direct a computer or other programmable data processing apparatus to operate in a particular manner, such that the instructions stored in the computer-readable memory produce an article of manufacture including instruction means which implement the acts specified in the flow chart and/or block diagram block or blocks. The computer program instructions may also be loaded to a computing device or other programmable data processing apparatus to cause operations to be performed on the computing device or other programmable apparatus to produce a computer implemented process such that the instructions which execute on the computing device or other programmable apparatus provide steps for implementing the acts specified in the flow chart and/or block diagram block or blocks.

(543) Any patents and applications and other references noted above, including any that may be listed in accompanying filing papers, are incorporated herein by reference. Aspects of the invention can be modified, if necessary, to employ the systems, functions, and concepts of the various references described above to provide yet further implementations of the invention. These and other changes can be made to the invention in light of the above Detailed Description. While the above description describes certain examples of the invention, and describes the best mode contemplated, no matter how detailed the above appears in text, the invention can be practiced in many ways. Details of the system may vary considerably in its specific implementation, while still being encompassed by the invention disclosed herein. As noted above, particular terminology used when describing certain features or aspects of the invention should not be taken to imply that the terminology is being redefined herein to be restricted to any specific characteristics, features, or aspects of the invention with which that terminology is associated. In general, the terms used in the following claims should not be construed to limit the invention to the specific examples disclosed in the specification, unless the above Detailed Description section explicitly defines such terms. Accordingly, the actual scope of the invention encompasses not only the disclosed examples, but also all equivalent ways of practicing or implementing the invention under the claims.

(544) To reduce the number of claims, certain aspects of the invention are presented below in certain claim forms, but the applicant contemplates other aspects of the invention in any number of claim forms. For example, while only one aspect of the invention is recited as a means-plus-function claim under 35 U.S.C. sec. 112(f) (AIA), other aspects may likewise be embodied as a means-plus-function claim, or in other forms, such as being embodied in a computer-readable medium. Any claims intended to be treated under 35 U.S.C. § 112(f) will begin with the words “means for,” but use of the term “for” in any other context is not intended to invoke treatment under 35 U.S.C. § 112(f). Accordingly, the applicant reserves the right to pursue additional claims after filing this application, in either this application or in a continuing application.

## Claims

1. A method comprising: receiving, by one or more processors, blockchain data from a plurality of nodes of a distributed ledger system; identifying a plurality of processing schema to connectors, wherein the plurality of processing schema includes a first processing schema and a second processing schema to ingest and process blockchain data generated by the plurality of nodes of the distributed ledger system, wherein the first processing schema indicates at least a first node of the

distributed ledger system to ingest the blockchain data and a first manner of processing the blockchain data, and wherein the second processing schema indicates at least a second node of the distributed ledger system to ingest the blockchain data and a second manner of processing the blockchain data, wherein the first manner includes at least a first filter or enrichment parameter and the second manner includes at least a second filter or enrichment parameter; processing the blockchain data according to the first processing schema, wherein the first processing schema is assigned to a first connector to automatically ingest the blockchain data from at least the first node, process the blockchain data according to the first processing schema, and generate first processed blockchain data, wherein the first processing schema is dynamically assigned to the first connector; processing the blockchain data according to the second processing schema, wherein the second processing schema is assigned to a second connector to automatically ingest the blockchain data from at least the second node, process the blockchain data according to the second processing schema, and generate second processed blockchain data, wherein the first connector and the second connector are separate from the distributed ledger system, wherein the second processing schema is dynamically assigned to the second connector; routing the first processed blockchain data and the second processed blockchain data to at least one component of a data intake and query system; filtering the first processed blockchain data and the second processed blockchain data to remove duplicated data; generating a first virtual index associated with the first processed blockchain data and generating a second virtual index associated with the second processed blockchain data; implementing an External Result Provider (ERP) process to provide an interface through which a search head to access the first virtual index and second virtual index for execution of a query.

2. The method of claim 1, wherein at least a portion of the blockchain data comprises transaction data for a blockchain transaction between a first node of the plurality of nodes and a second node of the plurality of nodes.

3. The method of claim 1, wherein the first connector and the second connector ingest a same portion of the blockchain data from a same node of the plurality of nodes, process the same portion of the blockchain data according to different processing schemas, and generate different processed blockchain data.

4. The method of claim 1, wherein the blockchain data is associated with a blockchain network, wherein the first connector and the second connector process the blockchain data based at least in part on one or more characteristics of the blockchain network.

5. The method of claim 1, wherein at least one of the plurality of nodes is internal to the data intake and query system.

6. The method of claim 1, wherein the blockchain data is associated with a first blockchain network, the method further comprising: receiving, by the one or more processors, additional blockchain data generated by an additional plurality of nodes, the additional blockchain data associated with a second blockchain network; implementing a third connector to ingest the additional blockchain data, process the additional blockchain data according to the first processing schema, and generate third processed blockchain data; implementing a fourth connector to ingest the additional blockchain data, process the additional blockchain data according to the second processing schema, and generate fourth processed blockchain data; and routing the third processed blockchain data and the fourth processed blockchain data to the data intake and query system.

7. The method of claim 1, wherein the blockchain data is associated with a first blockchain network, the method further comprising: receiving, by the one or more processors, additional blockchain data generated by an additional plurality of nodes, the additional blockchain data associated with a second blockchain network; implementing a third connector to ingest the additional blockchain data, process the additional blockchain data according to a third processing schema, and generate third processed blockchain data; implementing a fourth connector to ingest the additional blockchain data, process the additional blockchain data according to a fourth processing schema, and generate fourth processed blockchain data; and routing the third processed

blockchain data and the fourth processed blockchain data to the data intake and query system.

8. The method of claim 1, wherein the blockchain data is associated with a first blockchain network, the method further comprising: receiving, by the one or more processors, additional blockchain data generated by an additional plurality of nodes, the additional blockchain data associated with a second blockchain network; implementing a third connector to ingest the additional blockchain data, process the additional blockchain data according to a third processing schema, and generate third processed blockchain data; implementing a fourth connector to ingest the additional blockchain data, process the additional blockchain data according to a fourth processing schema, and generate fourth processed blockchain data; at least one of: correlating at least one of the third processed blockchain data or the fourth processed blockchain data to at least one of the first processed blockchain data or the second processed blockchain data, or combining at least one of the third processed blockchain data or the fourth processed blockchain data and at least one of the first processed blockchain data or the second processed blockchain data; and routing the third processed blockchain data and the fourth processed blockchain data to the data intake and query system.

9. The method of claim 1, further comprising: receiving, by the one or more processors, non-blockchain data generated by a computing device; processing the non-blockchain data to generate processed non-blockchain data; and routing the processed non-blockchain data to the data intake and query system.

10. The method of claim 1, further comprising: receiving, by the one or more processors, non-blockchain data generated by a computing device; processing the non-blockchain data to generate processed non-blockchain data; at least one of: correlating at least one of the first processed blockchain data or the second processed blockchain data to the processed non-blockchain data, or combining at least one of the first processed blockchain data or the second processed blockchain data and the processed non-blockchain data; and routing the processed non-blockchain data to the data intake and query system.

11. The method of claim 1, further comprising filtering the first processed blockchain data and the second processed blockchain data to remove duplicated data.

12. The method of claim 1, further comprising updating the first processing schema to generate an updated first processing schema, wherein additional blockchain data is processed according to the updated first processing schema based at least in part on updating the first processing schema.

13. The method of claim 1, further comprising removing the first processing schema, wherein additional blockchain data is not processed according to the first processing schema based at least in part on removing the first processing schema.

14. The method of claim 1, further comprising adding a third processing schema, wherein additional blockchain data is processed according to the first processing schema, the second processing schema, and the third processing schema based at least in part on adding the third processing schema.

15. The method of claim 1, wherein the blockchain data is associated with a blockchain network, further comprising: causing display, via a display device, of a user interface comprising one or more visualizations associated with the blockchain network, the one or more visualizations generated based at least in part on the first processed blockchain data and the second processed blockchain data.

16. The method of claim 1, wherein the blockchain data is associated with a blockchain network, further comprising: causing display, via a display device, of a user interface comprising one or more visualizations associated with the blockchain network, the one or more visualizations generated based at least in part on the first processed blockchain data and the second processed blockchain data, wherein the one or more visualizations indicate a blockchain status and operational device status of at least one of the plurality of nodes.

17. The method of claim 1, wherein the first processing schema is assigned to the first connector

based on observed or monitored query characteristics.

18. A computing system of a data ingestion system, the computing system comprising: a data store including computer-executable instructions; and one or more processors configured to execute the computer-executable instructions, wherein execution of the computer-executable instructions causes the computing system to: receive blockchain data from a plurality of nodes of a distributed ledger system; identify a plurality of processing schema to connectors, wherein the plurality of processing schema includes a first processing schema and a second processing schema to ingest and process blockchain data generated by the plurality of nodes of the distributed ledger system, wherein the first processing schema indicates at least a first node of the distributed ledger system to ingest the blockchain data and a first manner of processing the blockchain data, and wherein the second processing schema indicates at least a second node of the distributed ledger system to ingest the blockchain data and a second manner of processing the blockchain data, wherein the first manner includes at least a first filter or enrichment parameter and the second manner includes at least a second filter or enrichment parameter; process the blockchain data according to the first processing schema, wherein the first processing schema is assigned to a first connector to automatically ingest the blockchain data from at least the first node, process the blockchain data according to the first processing schema, and generate first processed blockchain data, wherein the first processing schema is dynamically assigned to the first connector; processing the blockchain data according to the second processing schema, wherein the second processing schema is assigned to a second connector to automatically ingest the blockchain data from at least the second node, process the blockchain data according to the second processing schema, and generate second processed blockchain data, wherein the first connector and the second connector are separate from the distributed ledger system, wherein the second processing schema is dynamically assigned to the second connector; route the first processed blockchain data and the second processed blockchain data to at least one component of a data intake and query system; filter the first processed blockchain data and the second processed blockchain data to remove duplicated data; generate a first virtual index associated with the first processed blockchain data and generating a second virtual index associated with the second processed blockchain data; implement an External Result Provider (ERP) process to provide an interface through which a search head to access the first virtual index and second virtual index for execution of a query.

19. The system of claim 18, wherein at least one of the plurality of nodes is internal to the data intake and query system.

20. Non-transitory computer readable media comprising computer-executable instructions that, when executed by a computing system of a data ingestion system, cause the computing system to: receive blockchain data from a plurality of nodes of a distributed ledger system; identify a plurality of processing schema to connectors, wherein the plurality of processing schema includes a first processing schema and a second processing schema to ingest and process blockchain data generated by the plurality of nodes of the distributed ledger system, wherein the first processing schema indicates at least a first node of the distributed ledger system to ingest the blockchain data and a first manner of processing the blockchain data, and wherein the second processing schema indicates at least a second node of the distributed ledger system to ingest the blockchain data and a second manner of processing the blockchain data, wherein the first manner includes at least a first filter or enrichment parameter and the second manner includes at least a second filter or enrichment parameter; process the blockchain data according to the first processing schema, wherein the first processing schema is assigned to a first connector to automatically ingest the blockchain data from at least the first node, process the blockchain data according to the first processing schema, and generate first processed blockchain data, wherein the first processing schema is dynamically assigned to the first connector; processing the blockchain data according to the second processing schema, wherein the second processing schema is assigned to a second connector to automatically ingest the blockchain data from at least the second node, process the blockchain data according to

the second processing schema, and generate second processed blockchain data, wherein the first connector and the second connector are separate from the distributed ledger system, wherein the second processing schema is dynamically assigned to the second connector; and route the first processed blockchain data and the second processed blockchain data to at least one component of a data intake and query system; filter the first processed blockchain data and the second processed blockchain data to remove duplicated data; generate a first virtual index associated with the first processed blockchain data and generating a second virtual index associated with the second processed blockchain data; implement an External Result Provider (ERP) process to provide an interface through which a search head to access the first virtual index and second virtual index for execution of a query.

---