

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12400068
Kind Code	B2
Date of Patent	August 26, 2025
Inventor(s)	Gotmare; Akhilesh Deepak et al.

Systems and methods for natural language code search

Abstract

Embodiments are directed to translating a natural language query into a code snippet in a programming language that semantically represents the query. The embodiments include a cascading neural network that includes an encoder network and a classifier network. The encoder network being faster but less accurate than the classifier network. The encoder network is trained using a contrastive learning framework to identify code candidates from a large set of code snippets. The classifier network is trained using a binary classifier to identify the code snippet that semantically represents the query from the code candidates.

Inventors: Gotmare; Akhilesh Deepak (Singapore, SG), Li; Junnan (Singapore, SG), Hoi; Chu Hong (Singapore, SG)

Applicant: Salesforce, Inc. (San Francisco, CA)

Family ID: 1000008782334

Assignee: Salesforce, Inc. (San Francisco, CA)

Appl. No.: 17/587984

Filed: January 28, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20230109681 A1	Apr. 13, 2023

Related U.S. Application Data

us-provisional-application US 63252393 20211005

Publication Classification

Int. Cl.: G06F40/151 (20200101); G06F40/30 (20200101); G06F40/40 (20200101); G06N3/04 (20230101)

U.S. Cl.:

CPC G06F40/151 (20200101); G06F40/30 (20200101); G06F40/40 (20200101); G06N3/04 (20130101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
11132512	12/2020	Wu	N/A	G06N 3/044
11734339	12/2022	Kale	707/722	G06F 18/21
2012/0254143	12/2011	Varma	707/706	G06F 40/30
2020/0117446	12/2019	Smith	N/A	G06F 8/71
2020/0334520	12/2019	Chen	N/A	G06F 40/20
2021/0191696	12/2020	Ibarra Von Borstel	N/A	G06F 8/36
2021/0303989	12/2020	Bird	N/A	G06F 16/95
2023/0042051	12/2022	Clement	N/A	G06N 3/091
2023/0100376	12/2022	Liu	704/9	G06F 40/30

OTHER PUBLICATIONS

CoSQA: 20,000+ Web Queries for Code Search and Question Answering, arXiv:2105.13239v1 (Year: 2021). cited by examiner

Lobato, "Neural networks with optimal accuracy and speed in their predictions," Toward Data Science.com, Jul. 20, 2017 (available at: <https://towardsdatascience.com/neural-networks-with-optimal-accuracy-and-speed-in-their-predictions-d2cdc3b21b50>) (Year: 2017). cited by examiner
Pengcheng Yin, Bowen Deng, Edgar Chen, Bogdan Vasilescu, Graham Neubig, "Learning to Mine Aligned Code and Natural Language Pairs from Stack Overflow", 2018, IEEE, ACM/IEEE 15th International Conference on Mining Software Repositories, pp. 476-486 (Year: 2018). cited by examiner

International Search Report and Written Opinion for PCT/US2022/077458, dated Jan. 4, 2023, 11 pages. cited by applicant

Gotmare et al., "Cascaded Fast and Slow Models for Efficient Semantic Code Search", arXiv.Org., Cornell University Library, 201 Olin Library Cornell University Ithaca, NY 14853, Oct. 15, 2021, XP091077176, pp. 1-12. cited by applicant

Guo et al., "GraphCodeBert: Pre-Training Code Representations with Data Flow", Published as conference paper at ICLR 2021, Sep. 13, 2021, Retrieved from the Internet:

URL:<https://arxiv.org/pdf/2009.08366.pdf>, pp. 1-18. cited by applicant

Salza et al., "On the Effectiveness of Transfer Learning for Code Search", arXiv.Org., 201 Olin Library Cornell University Ithaca, NY 14853, Aug. 12, 2021, XP091032288, 16 pages. cited by applicant

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a nonprovisional of and claims priority under 35 U.S.C. 119 to U.S. Provisional Application No. 63/252,393, filed Oct. 5, 2021, which is hereby expressly incorporated by reference herein in its entirety.

TECHNICAL FIELD

(1) The embodiments relate generally to machine learning systems and natural language processing (NLP), and more specifically to searching code snippets using natural language.

BACKGROUND

(2) Artificial intelligence (AI) models have been widely used in a variety of applications. Some AI models may be used to search and/or generate code snippets in programming languages in response to a natural language input. For example, a natural language input may describe a function such as “filter the sales records that occurred at the zip code 94070,” and the AI model may generate or search a code segment (e.g., in Python, C#, etc.) that implements this function. Existing code generation systems have focused on either improving the speed of natural language search or on improving the accuracy of the natural language search. However, these existing natural language search methods largely struggle with a tradeoff between efficiency and exhaustiveness of the search.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1 is a simplified diagram of a computing device that implements a code generator, according to some embodiments.

(2) FIG. 2 is a simplified diagram of a code generator, according to some embodiments.

(3) FIG. 3 is a simplified diagram of a method for training the code generator, according to some embodiments.

(4) FIG. 4 is a simplified diagram of a method for determining a code snippet that is semantically equivalent to a natural language query, according to some embodiments.

(5) In the figures, elements having the same designations have the same or similar functions.

DETAILED DESCRIPTION

(6) Natural language queries are being used to improve search in different areas, e.g., web search, database search, legal search, and/or the like. There is also interest in searching large sets of code snippets using natural language queries. Organizations that have large code repositories may benefit from indexing and searching through the code and reuse code that is known to function properly. Some recent approaches to natural language search of code and code snippets leverage pairs of natural language and source code sequences to train a text-to-code search model to search a sample of code snippets.

(7) One approach to train the model includes using a contrastive learning framework. The model may be a fast encoder neural network, also referred to as a fast encoder. In the contrastive learning framework, pairs of natural language and program language sequences that semantically match are pulled together while pairs that do not semantically match are pushed apart. Fast encoder networks may use contrastive learning. Fast encoders networks may be efficient for scenarios that include searching a large number of candidate code snippets at the expense of accuracy in semantics

matching.

(8) Another approach to train a model uses a binary classifier. This type of a model uses a trained binary classifier that receives a natural language and programming language sequences as inputs and predicts whether the natural language and programming language sequences match semantically. Models using binary classifiers may be considered slow classifiers. Slow classifiers, while more accurate, become infeasible when searching a large number of candidate code snippets due to an amount of time the models take to analyze the code snippets against the natural language sequence. In other words, models trained using contrastive learning framework may be faster by at least a factor of ten, but may also be less accurate by at least a factor of ten or more than models that use a binary classifier.

(9) To improve natural language searching of large numbers of code snippets, embodiments are directed to a cascading neural network model that includes both fast encoder model and an accurate classifier model. The cascading neural network model provides improved natural language search efficiency of large sets of code snippets. Specifically, the cascading neural network model is a hybrid approach that combines a fast encoder network and a slow classifier network. First, the encoder network determines a top K code candidates from the set of code snippets based on a natural language query. Second, the top K code candidates are passed through a slow classifier network that pairs each of the code candidates with the natural language query and generates confidence scores for each pair. The code snippet with a highest confidence score may be the code snippet that semantically matches the natural language query.

(10) The number K may denote a threshold that identifies the number of code candidates that the encoder network may generate. The K threshold is preferably much smaller than the size of the set of code snippets. If the K threshold is too small there is an increased likelihood of missing the correct code snippet and if K threshold is too large it may be infeasible to efficiently run the second stage slow classifier.

(11) In some embodiments, memory overhead for storing a fast encoder network and a slow classifier network may be minimized by sharing or partially sharing the weights of the network. For example, a transformer encoder of the fast encoder and slow classifier may be shared by training the transformer encoder to be used by both the fast encoder network and the slow classifier network.

(12) As used herein, the term “network” may comprise any hardware or software-based framework that includes any artificial intelligence network or system, neural network or system and/or any training or learning models implemented thereon or therewith.

(13) As used herein, the term “module” may comprise hardware or software-based framework that performs one or more functions. In some embodiments, the module may be implemented on one or more neural networks.

(14) FIG. 1 is a simplified diagram of a computing device that implements a code generator, according to some embodiments described herein. As shown in FIG. 1, computing device **100** includes a processor **110** coupled to memory **120**. Operation of computing device **100** is controlled by processor **110**. And although computing device **100** is shown with only one processor **110**, it is understood that processor **110** may be representative of one or more central processing units, multi-core processors, microprocessors, microcontrollers, digital signal processors, field programmable gate arrays (FPGAs), application specific integrated circuits (ASICs), graphics processing units (GPUs) and/or the like in computing device **100**. Computing device **100** may be implemented as a stand-alone subsystem, as a board added to a computing device, and/or as a virtual machine.

(15) Memory **120** may be used to store software executed by computing device **100** and/or one or more data structures used during operation of computing device **100**. Memory **120** may include one or more types of machine readable media. Some common forms of machine readable media may include floppy disk, flexible disk, hard disk, magnetic tape, any other magnetic medium, CD-ROM, any other optical medium, punch cards, paper tape, any other physical medium with patterns of

holes, RAM, PROM, EPROM, FLASH-EPROM, any other memory chip or cartridge, and/or any other medium from which a processor or computer is adapted to read.

(16) Processor **110** and/or memory **120** may be arranged in any suitable physical arrangement. In some embodiments, processor **110** and/or memory **120** may be implemented on a same board, in a same package (e.g., system-in-package), on a same chip (e.g., system-on-chip), and/or the like. In some embodiments, processor **110** and/or memory **120** may include distributed, virtualized, and/or containerized computing resources, as well as multiple processors. Consistent with such embodiments, processor **110** and/or memory **120** may be located in one or more data centers and/or cloud computing facilities.

(17) In some examples, memory **120** may include non-transitory, tangible, machine readable media that includes executable code that when run by one or more processors (e.g., processor **110**) may cause the one or more processors to perform the methods described in further detail herein. For example, as shown, memory **120** includes instructions for a natural language (NL) processing module, such as a code generator **130** that may be used to implement and/or emulate the systems and models, and/or to implement any of the methods described further herein. In some examples, the code generator **130**, may receive an input **140**, e.g., such as natural language text or query or computer code, via a data interface **115**. The data interface **115** may be any of a user interface that receives input **140** from a user, or a communication interface that receives or retrieves input **140** stored in memory **120** or another memory storage such as a database. The code generator **130** may generate an output **150**, such as a programmable language (PL) sequence, code, or code snippet that is semantically equivalent to the natural language text or query. In some embodiments, code generator **130** may include a cascading neural network that includes an encoder network **132** and a classifier network **134**, such that the output of encoder network **132** may be input, in part, into classifier network **134**.

(18) FIG. 2 is a block diagram **200** of a code generator, according to some embodiments. As illustrated in FIG. 2, the code generator **130** includes encoder network **132** and classifier network **134**. The code generator **130** receives a natural language query **202** in a natural language query or text. Natural language query **202** may be input **140** discussed in FIG. 1. The natural language query **202** may be a human written or spoken text, such as “filter the sales records that occurred at the zip code 94070” that code generator **130** may translate into a programming language sequence, such as a code snippet. The code generator **130** passes natural language query **202** through encoder network **132**. Encoder network **132** may generate a K number of code candidates **204A-K**. The code candidates **204A-K** may be code snippets in a programming language that may semantically represent and/or semantically match the natural language query **202**. The classifier network **134** may receive pairs of code candidates **204A-K** the natural language query **202**. Each pair in the pairs may include one of candidates **204A-K** and natural language query **202**. Classifier network **134** may generate a code snippet **206** that is a semantic representation of the natural language query **202**.

(19) In some embodiments, encoder network **132** is substantially faster, e.g. at least by a factor of ten or more, than the classifier network **134**. In fact, due to speed of the encoder network **132**, encoder network **132** may quickly determine code candidates **204A-K** from a large set of available code snippets. The classifier network **134**, on the other hand, is slower than encoder network **132**, but is substantially more accurately, e.g., at least by at least a factor of ten or more, in identifying a code snippet that semantically matches the natural language query **202**. As illustrated in FIG. 2, classifier network **134** receives pairs of code candidates **204A-K** and natural language query **202** and identifies code snippet **206** that is a semantic representation of the natural language query **202**. By using the hybrid approach that includes encoder network **132** and classifier network **134**, code generator **130** improves the speed and accuracy for determining the code snippet **206** that is a semantic representation of the natural language query **202**.

(20) In some embodiments, encoder network **132** may be or include a bidirectional encoder

representations from transformers (BERT) network or a variant of a BERT network. The BERT network or a variant of the BERT network may be pre-trained on programming language sequences in a variety of programming languages to retrieve code snippets from text input. Example pre-trained BERT networks may be GraphCodeBERT or CodeBERT. An example programming languages may be Ruby, JavaScript, Go, Python, Java, C, C++, C#, Php, or the like. During a training phase, to recognize code candidates **204**, encoder network **132** may further be trained on a the contrastive learning framework using a bimodal dataset. In the bimodal dataset, the positive pairs, that are representations of natural language queries and programming language sequences that match in semantics are pulled together. On the other hand, representations of negative pairs, that are randomly paired natural language queries and programming language sequences, are pushed apart. A contrastive loss function, such as the infoNCE loss function may be used to train encoder network **132**, and is replicated below:

$$(21) \quad \mathcal{L}_{\text{infoNCE}} = \frac{1}{N} \cdot \text{Math.} \cdot \log \frac{\exp(f(x_i)^T f(y_i))}{\sum_{j \in B} \exp(f(x_i)^T f(y_j))} \quad (1)$$

where $f(\text{sub.}\theta(x.\text{sub.}i))$ is the dense representation for the natural language input $x.\text{sub.}i$, $y.\text{sub.}i$ is the corresponding semantically equivalent programming language sequence, N is the number of training examples in the bimodal dataset, σ is a temperature hyper-parameter, and B denotes the current training minibatch. The encoder network **132** may be trained until the contrastive loss function is minimized.

(22) Once trained, encoder network **132** may receive a set of candidate code snippets $C = \{\text{custom character.sub.1}, \text{custom character.sub.2}, \dots, \text{custom character.sub.}|C|\}$, shown as code snippets **208**. Code snippets **208** may include potential code snippets, universe of code snippets, available code snippets., etc., that may correspond to various natural language queries. Code snippets **208** may be encoded into an index $\{f(\text{sub.}\theta(\text{custom character.sub.}j)) \forall \text{sub.}j \in C\}$, shown as code snippet index **210**. Code snippet index **210** may be an index of the encodings of each code snippet in code snippets **208**. Encoder network **132** may encode the set of code snippets offline, e.g., prior to receiving natural language query **202** from which code generator **130** determines code snippet **206**. Code snippet index **210** may be stored within encoder network **132** or elsewhere in memory **120** described in FIG. 1.

(23) In some embodiments, after generating the code snippet index **210**, encoder network **132** may receive a natural query $x.\text{sub.}i$ (natural language query **202**), compute $f(\text{sub.}\theta(x.\text{sub.}i))$, query the code snippet index **210**, and return the code snippet from C (code snippets **208**) corresponding to the nearest neighbor or neighbors in the code snippet index **210**. The neighbor(s) may be computed using a distance metric determined by a similarity function, e.g., a cosine similarity function. The rank r_i assigned to the correct code snippet from the set of code snippets C (code snippets **208**) for the natural language query $x.\text{sub.}i$ may then be used to compute the mean reciprocal ranking (MRR) metric

$$(24) \quad \frac{1}{N_{\text{test}}} \cdot \text{Math.} \cdot \sum_{i=1}^{N_{\text{test}}} \frac{1}{r_i}.$$

From the MRR metric, the code candidates **204** with a rank included in the MRR or a certain distance from the rank in the MRR may be determined. In some embodiments, the number of code candidates **204** may be governed by a hyperparameter K which may be a threshold values that cause encoder network **132** to identify the top K candidates, e.g. code candidates **204A-K**.

(25) In some embodiments, classifier network **134** may also be or include a bidirectional encoder representations from transformers (BERT) network or a variant of BERT network. The BERT network or a variant of the BERT network may be pre-trained on programming language sequences to retrieve code snippets from text input. Example pre-trained BERT networks may be GraphCodeBERT or CodeBERT and an example programming language languages may be Ruby, JavaScript, Go, Python, Java, C, C++, C#, Php, or the like.

(26) Classifier network **134** may receive natural language query $x.\text{sub.}i$ (shown as **202**) and a

programming language sequence custom character.sub.i (one of code candidates **204A-K** or another code sequence) as input, encode the natural language input x.sub.i and the code sequence jointly and perform binary classification. The binary classification may predict whether the natural language input x.sub.i and a code sequence custom character.sub.i match or does not match in semantics. In some embodiments, classifier network **134** may receive a concatenation of the natural language input x.sub.i and a code sequence custom character.sub.j, such as [x.sub.i; custom character.sub.j].

(27) Classifier network **134** may be trained for binary classification using training batches. The training batches may include pairs where each pair includes a natural language query and a code snippet. The training batches may be batches for a bimodal dataset where positive pairs denote semantic matches between the natural language query and the code snippet and the negative pairs denote semantic mismatches. Given a set of pairs that include natural language query and a semantically programming language sequence {x.sub.i, custom character.sub.i}.sub.i=1.sup.N, the cross-entropy objective function for this training scheme may be:

$$(28) \mathcal{L}_{CE} = -\frac{1}{N} \sum_{i=1, j \neq i}^N \text{Math.} \log p(x_i, i) + \log(1 - p(x_i, j)) \quad (2)$$

where $p.\text{sub}.\theta(x.\text{sub}.i, \text{img alt="code snippet icon" data-bbox="254 324 278 338"/>custom character.\text{sub}.j)$ represents the probability that the natural language sequence x.sub.i semantically matches the programming language sequence custom character .sub.j, as predicted by the classifier. Classifier network **134** may be trained until the cross-entropy objective function is minimized.

(29) From a training minibatch of positive pairs {x.sub.i, custom character.sub.i} $\forall i \in B$, a training batch of negative pairs may be generated. For example, a negative pair may be generated by randomly selecting a programming language sequence custom character.sub.j ($j \in B; j \neq i$) from the programming language sequences in the minibatch and pairing the selected sequence with x.sub.i. When classifier network **134** includes a transformer encoder based classifier, the interactions between the natural language tokens and programming language tokens in the self-attention layers may help in improving the precision of the classifier network **134**.

(30) Once trained, classifier network **134** may determine code snippet **206** from natural language query **202** and code candidates **204**. For example, during inference, classifier network **134** may receive multiple pairs as inputs, each pair including the natural language sequence x.sub.i (e.g. natural language query **202**) with a code snippet custom character.sub.j (one of code candidates **204A-K**) from a set of candidate code snippets $C = \{\text{img alt="code snippet icon" data-bbox="511 625 535 639"/>custom character.\text{sub}.1, \text{img alt="code snippet icon" data-bbox="541 625 565 639"/>custom character.\text{sub}.2, \dots, \text{img alt="code snippet icon" data-bbox="571 645 595 659"/>custom character.\text{sub}.|C|\}$ (code candidates **204A-K**). Classifier network **134** may generate a confidence score for each pair and rank each code candidate in the code candidate **204A-K** according to the confidence scores. The confidence score may be a probability having a measure, e.g., from zero to one, with the values closer to one indicating a greater probability of a match, and the values closer to zero indicating a greater probability of a mismatch. The code snippet custom character.sub.j (a code candidate in the code candidates **204A-K**) that corresponds to a pair with the highest score may be a semantic match with the natural language sequence x.sub.i (natural language query **202**).

(31) As discussed above, the code generator **130** discussed herein includes a cascade of networks, such as encoder network **132** and classifier network **134** which combines the speed of the fast encoder network **132** with precision of the classifier network **134** in a two stage process. In the first stage, encoder network **132** receives natural language query **202** and generates code candidates **204A-K** from the set of code snippets C (code snippets **208**). Encoder network **132** may determine encodings of the natural language query **202** and match the encodings against the code snippet index **210** of the code snippet **208** using a distance function. In some embodiments, encoder network **132** may determine a K number of code candidates **204A-K** where K is a configurable candidate threshold that may be a hyperparameter. Typically, the K number of candidates are the

top candidates that have the closest distance in the code snippet index **210** to the encodings of the natural language query **202**.

(32) In the second stage, the code candidates **204** are paired with the natural language query **202**. Example pairs may be **202-204A**, **202-204B**, . . . , **202-204K**. Classifier network **134** receives pairs **202-204A**, **202-204B**, . . . , **202-204K**. For each pair in pairs **202-204A**, **202-204B**, . . . , **202-204K**, classifier network **134** returns a confidence score that the natural language query **202** semantically matches a corresponding one of code candidates **204A-K** using a binary classifier. Based on the confidence scores associated with pairs **202-204A**, **202-204B**, . . . , **202-204K**, classifier network **134** selects the code snippet **206** that semantically matches the natural language query **202**. In some instances, code snippet **206** may correspond to a pair having a highest confidence score.

(33) As discussed above, encoder network **132**, while computationally faster is also less accurate, in determining a code snippet that semantically matches natural language query, than classifier network **134**. In a scheme where $K \ll |C|$, adding classifier network **134** in sequence with encoder network **132** may add a minor computational overhead. The second stage where classifier network **134** refines code candidates **204A-K** improves the retrieval performance provided that the value of K is set such that the recall of the encoder network **132** is reasonably high. In some embodiments, K may be a hyperparameter. Setting a very low K would lead to high likelihood of missing the code snippet **206** in the set of code candidates **204** passed to classifier network **134**. On the other hand, setting a high K would make the scheme infeasible for retrieval with the classifier network **134**. However, setting K to a value, such as ten, already offers significant gains in retrieval performance over conventional code generation systems, which only marginal gains when K is set to 100 and above.

(34) In some embodiments, encoder network **132** and classifier network **134** may share a portion of a neural network structure. For example, encoder network **132** and classifier network **134** may share weights of the layers in the transformer encoder in the BERT network. Sharing the neural network structure minimizes a memory overhead incurred by encoder network **132** and classifier network **134**. Sharing the neural network structure, e.g. transformer layers, by encoder network **132** and classifier network **134** may be achieved by training the transformer encoder with a joint objective of infoNCE (custom character.sub.infoNCE) shown in Eq. (1) and binary cross-entropy (custom character.sub.BCE) shown in Eq. (2). While the number of parameters in this shared variant would be nearly half of when the transformer layers are not shared, the computational cost during the inference may be similar or the same.

(35) In the shared embodiment, the classifier network **134** may have an additional classification layer or head that determines confidence scores for the pairs **202-204A**, **202-204B**, . . . , **202-204K**. Classifier network **134** would include the classification head on top of the transformer encoder. Further, the shared neural network structure may receive three inputs, natural language query **202**, set of candidate code snippets C (code snippets **208**), and pairs **202-204A**, **202-204B**, . . . , **202-204K**. In the shared embodiment, two passes are made through the shared layers of the network, where the natural language query **202** is an input during a first pass, and pairs **202-204A**, **202-204B**, . . . , **202-204K** are input during a second pass.

(36) FIG. 3 is a simplified diagram of a method **300** for training a code generator, according to some embodiments. One or more of the processes **302-304** of method **300** may be implemented, at least in part, in the form of executable code stored on non-transitory, tangible, machine-readable media that when run by one or more processors may cause the one or more processors to perform one or more of the processes **302-304**.

(37) At process **302**, an encoder network is trained. For example, encoder network **132** that may be a pre-trained BERT network may further be trained on a contrastive learning framework to identify code snippets that semantically matches the natural language sequences. The contrastive loss function used to train encoder network **132** may be a contrastive loss function, such as infoNCE loss function. The training may include batches of negative and positive pairs, each pair including a

natural language sequence and a programming language sequence. The training may continue iteratively until the infoNCE loss function is minimized.

(38) At process **304**, a classifier network is trained. For example, classifier network **134** that may be a pre-trained BERT network may be trained on binary classification to determine a probability score that the code snippets match the natural language sequences. The cross-entropy objective function may be used to train classifier network **134**. The training may include batches of negative and positive pairs, each pair including a natural language sequence and a programming language sequence. The training may continue iteratively until the cross-entropy objective function is minimized.

(39) FIG. **4** is a simplified diagram of a method **400** for generating a code snippet that is semantically equivalent to a natural language query, according to some embodiments. One or more of the processes **402-408** of method **400** may be implemented, at least in part, in the form of executable code stored on non-transitory, tangible, machine-readable media that when run by one or more processors may cause the one or more processors to perform one or more of the processes **402-408**.

(40) At process **402**, a code snippet index is generated. For example, encoder network **132** receives code snippets **208** that may semantically correspond to numerous natural language queries. Encoder network **132** encodes the code snippets **208** and generates a code snippet index **210** that corresponds to the encoded code snippets. Process **402** may occur after encoder network **132** is trained and prior to encoder network **132** processing natural language query **202**.

(41) At process **404**, code candidates for a natural language query are generated. For example, encoder network **132** may receive natural language query **202** and generate encodings for the natural language query **202**. Encoder network **132** may use the code snippet index **210** to match the encodings of the natural language query **202** to encodings of code snippets **208** to identify code candidates **204A-K** that may semantically match natural language query **202**. As discussed above, the number of code candidates **204A-K** may be set using a number **K** which may be a hyperparameter.

(42) At process **406**, pairs that include natural language query and code candidates are generated. For example, code generator **130** may generate pairs **202-204A**, **202-204B**, . . . , **202-204K**, where each pair includes natural language query **202** and one of code candidates **204A-K**.

(43) At process **408**, a code snippet is determined. For example, classifier network **104** may receive the pairs **202-204A**, **202-204B**, . . . , **202-204K** and determine a confidence score for each pair. The pair with a highest confidence score may be code snippet **206** that semantically matches natural language query **202**.

(44) Some examples of computing devices, such as computing device **100** may include non-transitory, tangible, machine readable media that include executable code that when run by one or more processors (e.g., processor **110**) may cause the one or more processors to perform the processes of methods **300-400**. Some common forms of machine readable media that may include the processes of methods **300-400** are, for example, floppy disk, flexible disk, hard disk, magnetic tape, any other magnetic medium, CD-ROM, any other optical medium, punch cards, paper tape, any other physical medium with patterns of holes, RAM, PROM, EPROM, FLASH-EPROM, any other memory chip or cartridge, and/or any other medium from which a processor or computer is adapted to read.

(45) This description and the accompanying drawings that illustrate inventive aspects, embodiments, implementations, or applications should not be taken as limiting. Various mechanical, compositional, structural, electrical, and operational changes may be made without departing from the spirit and scope of this description and the claims. In some instances, well-known circuits, structures, or techniques have not been shown or described in detail in order not to obscure the embodiments of this disclosure. Like numbers in two or more figures represent the same or similar elements.

(46) In this description, specific details are set forth describing some embodiments consistent with the present disclosure. Numerous specific details are set forth in order to provide a thorough understanding of the embodiments. It will be apparent, however, to one skilled in the art that some embodiments may be practiced without some or all of these specific details. The specific embodiments disclosed herein are meant to be illustrative but not limiting. One skilled in the art may realize other elements that, although not specifically described here, are within the scope and the spirit of this disclosure. In addition, to avoid unnecessary repetition, one or more features shown and described in association with one embodiment may be incorporated into other embodiments unless specifically described otherwise or if the one or more features would make an embodiment non-functional.

(47) Although illustrative embodiments have been shown and described, a wide range of modification, change and substitution is contemplated in the foregoing disclosure and in some instances, some features of the embodiments may be employed without a corresponding use of other features. One of ordinary skill in the art would recognize many variations, alternatives, and modifications. Thus, the scope of the invention should be limited only by the following claims, and it is appropriate that the claims be construed broadly and in a manner consistent with the scope of the embodiments disclosed herein.

Claims

1. A method for translating a natural language query into a code snippet in a programming language, the method comprising: generating, at an encoder network, a plurality of code snippet indices and a natural language query representation from a plurality of code snippets and the natural language query, respectively; determining one or more code candidates based on distance metrics between the natural language query representations and the plurality of code snippet indices; generating text-code pairs from the natural language query and the determined one or more code candidates based on the distance metrics, a text-code pair including the natural language query and a code candidate from the determined one or more code candidates; jointly encoding, by a single encoder of a classifier network, an input sequence concatenating the natural language query and the code candidate from the text-code pair into one joint representation, wherein the single encoder and the encoder network share one or more encoder layers; predicting, by a classification head of the classifier network, a matching probability that the natural language query and the code candidate are a match using an input of the one joint representation outputted from the single encoder; and outputting the code candidate in the programming language in response to the natural language query based on the matching probability.
2. The method of claim 1, further comprising: training, the encoder network to determine the code candidates on a contrastive loss function.
3. The method of claim 1, further comprising: training, the classifier network to determine the code snippet from the pairs using a cross-entropy objective function.
4. The method of claim 1, wherein the classifier network sequentially follows the encoder network, and the encoder network is an order of magnitude faster and an order of magnitude less accurate than the classifier network.
5. The method of claim 1, wherein the encoder network is trained on a different loss function from the classifier network.
6. The method of claim 1, wherein the encoder network shares a portion of a neural network structure with the classifier network.
7. The method of claim 1, wherein generating the code candidates further comprises: generating encodings from the natural language query; and determining, using the code snippet index, the encodings of the code candidates that are within a distance determined by a distance function from the encodings of the natural language query.

8. The method of claim 1, wherein determining the code snippet further comprises: determining a confidence score that a code candidate of each pair is the semantic representation of the natural language query; ranking confidence scores of the pairs; and selecting a code candidate of a pair corresponding to a highest confidence score as the code snippet that is the semantic representation of the natural language query.

9. A system for translating a natural language query into a code snippet in a programming language, the system comprising: a memory configured to store a cascading neural network; a processor coupled to the memory and configured to execute instructions for causing the cascading neural network to: generate, at an encoder network of the cascading neural network, a plurality of code snippet indices and a natural language query representation from a plurality of code snippets and the natural language query, respectively; determine one or more code candidates based on distance metrics between the natural language query representation and the plurality of code snippet indices; generate text-code pairs from the natural language query and the determined one or more code candidates based on the distance metrics, a text-code pair including the natural language query and a code candidate from the determined one or more code candidates; jointly encode, by a single encoder of a classifier network, an input sequence concatenating the natural language query and the code candidate from the text-code pair into one joint representation, wherein the single encoder and the encoder network share one or more encoder layers; predict, by a classification head of the classifier network, a matching probability that the natural language query and the code candidate are a match using an input of the one joint representation outputted from the single encoder; and output the code candidate in the programming language in response to the natural language query, based on the matching probability.

10. The system of claim 9, wherein the processor is further configured to: train, the encoder network to determine the code candidates on a contrastive loss function; and train, the classifier network to determine the code snippet from the pairs using a cross-entropy objective function.

11. The system of claim 9, wherein the classifier network sequentially follows the encoder network, and the encoder network is an order of magnitude faster and an order of magnitude less accurate than the classifier network.

12. The system of claim 9, wherein the encoder network shares a portion of a neural network structure with the classifier network.

13. The system of claim 9, wherein to generate the code candidates the processor is further configured to: generate encodings from the natural language query; and determine, using the code snippet index, encodings of the code candidates that are within a distance determined by a distance function from the encodings of the natural language query.

14. The system of claim 9, wherein to determine the code snippet the processor is further configured to: determine a confidence score that a code candidate of each pair is a semantic representation of the natural language query; rank confidence scores of the pairs; and select a code candidate of a pair corresponding to a highest confidence score as the code snippet that is the semantic representation of the natural language query.

15. A non-transitory computer readable medium having instructions stored thereon, that when executed by a processor cause the processor to perform operations for translating a natural language query into a code snippet in a programming language, the operations comprising: generating, at an encoder network, a plurality of code snippet indices and a natural language query representation from a plurality of code snippets and the natural language query, respectively; determining one or more code candidates based on distance metrics between the natural language query representations and the plurality of code snippet indices; generating text-code pairs from the natural language query and the determined one or more code candidates based on the distance metrics, a text-code pair including the natural language query and a code candidate from the determined one or more code candidates; jointly encoding, by a single encoder of a classifier network, an input sequence concatenating the natural language query and the code candidate from

the text-code pair into one joint representation, wherein the single encoder and the encoder network share one or more encoder layers; predicting, by a classification head of the classifier network, a matching probability that the natural language query and the code candidate are a match using an input of the one joint representation outputted from the single encoder; and outputting the code candidate in the programming language in response to the natural language query based on the matching probability.

16. The non-transitory computer readable medium of claim 15, further comprising: training, the encoder network to determine the code candidates on a contrastive loss function; and training, the classifier network to determine the code snippet from the pairs using a cross-entropy objective function.

17. The non-transitory computer readable medium of claim 15, wherein the classifier network sequentially follows the encoder network, and the encoder network is an order of magnitude faster and an order of magnitude less accurate than the classifier network.

18. The non-transitory computer readable medium of claim 15, wherein the encoder network shares a portion of a neural network structure with the classifier network.

19. The non-transitory computer readable medium of claim 15, wherein generating the code candidates further comprises: generating encodings from the natural language query; and determining, using the code snippet index, encodings of the code candidates that are within a distance determined by a distance function from the encodings of the natural language query.

20. The non-transitory computer readable medium of claim 15, wherein determining the code snippet further comprises: determining a confidence score that a code candidate of each pair is a semantic representation of the natural language query; ranking confidence scores of the pairs; and selecting a code candidate of a pair corresponding to a highest confidence score as the code snippet that is the semantic representation of the natural language query.
