

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12400118
Kind Code	B2
Date of Patent	August 26, 2025
Inventor(s)	Keller; Alexander et al.

Incorporating a ternary matrix into a neural network

Abstract

Artificial neural networks (ANNs) are computing systems inspired by the human brain by learning to perform tasks by considering examples. These ANNs are typically created by connecting several layers of artificial neurons using connections, where each artificial neuron is connected to every other artificial neuron either directly or indirectly to create fully connected layers within the ANN. By substituting ternary matrices for one or more fully connected layers within the ANN, a complexity and resource usage of the ANN may be reduced, while improving the performance of the ANN.

Inventors: Keller; Alexander (Berlin, DE), Torcato Mordido; Gonalo Filipe (Potsdam, DE), Van Keirsbilck; Matthijs Jules (Berlin, DE)

Applicant: NVIDIA Corporation (Santa Clara, CA)

Family ID: 1000008778848

Assignee: NVIDIA CORPORATION (Santa Clara, CA)

Appl. No.: 17/356279

Filed: June 23, 2021

Prior Publication Data

Document Identifier	Publication Date
US 20220129755 A1	Apr. 28, 2022

Related U.S. Application Data

us-provisional-application US 63106287 20201027

Publication Classification

Int. Cl.: G06N3/082 (20230101); G06F7/50 (20060101); G06F7/523 (20060101); G06F7/58 (20060101)

U.S. Cl.:

CPC G06N3/082 (20130101); G06F7/50 (20130101); G06F7/523 (20130101); G06F7/582 (20130101);

Field of Classification Search

CPC: G06N (3/082); G06F (7/50); G06F (7/523); G06F (7/582)

References Cited

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
110909883	12/2019	CN	N/A
110942148	12/2019	CN	N/A

OTHER PUBLICATIONS

Boo, Yoonho, and Wonyong Sung. "Compression of Deep Neural Networks with Structured Sparse Ternary Coding." Journal of Signal Processing Systems 91 (2019): 1009-1019. (Year: 2019). cited by examiner

Gope, Dibakar, et al. "Ternary mobilenets via per-layer hybrid filter banks." Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops. 2020. (Year: 2020). cited by examiner

Hannun, Awni, et al. "Sequence-to-sequence speech recognition with time-depth separable convolutions." arXiv preprint arXiv: 1904.02619 (2019). (Year: 2019). cited by examiner

Heinrich, Mattias P., Max Blendowski, and Ozan Oktay. "TernaryNet: faster deep model inference without GPUs for medical 3D segmentation using sparse and binary convolutions." International journal of computer assisted radiology and surgery 13 (2018): 1311-1320. (Year: 2018). cited by examiner

Kundu, Abhisek, et al. "Ternary residual networks." arXiv preprint arXiv: 1707.04679 (2017). (Year: 2017). cited by examiner

Ling, Hefei, et al. "XwiseNet: action recognition with Xwise separable convolutions." Multimedia Tools and Applications 79 (2020): 26913-26926.

(Year: 2020). cited by examiner

Liu, Weiliang, et al. "Training an Artificial Neural Network with Op-amp Integrators Based Analog Circuits." 2018 IEEE CSAA Guidance, Navigation and Control Conference (CGNCC). IEEE, 2018. (Year: 2018). cited by examiner

Ambai, Mitsuru, et al. "Ternary weight decomposition and binary activation encoding for fast and compact neural network." (2016). (Year: 2016). cited by examiner

Prost-Boucle, Adrien, et al. "Scalable high-performance architecture for convolutional ternary neural networks on FPGA." 2017 27th International conference on field programmable logic and applications (Fpl). IEEE, 2017. (Year: 2017). cited by examiner

Qian, Zhongchao, et al. "Do we need fully connected output layers in convolutional networks?." arXiv preprint arXiv:2004.13587 (2020). (Year: 2020). cited by examiner

Rauf, Hafiz Tayyab, et al. "Training of artificial neural network using pso with novel initialization technique." 2018 international conference on innovation and intelligence for informatics, computing, and technologies (3ICT). IEEE, 2018. (Year: 2018). cited by examiner

Tsai, Chang-Hung, et al. "A 41.3/26.7 pJ per neuron weight RBM processor supporting on-chip learning/inference for IoT applications." IEEE Journal of Solid-State Circuits 52.10 (2017): 2601-2612. (Year: 2017). cited by examiner

Xu, Menglong, and Xiao-Lei Zhang. "Depthwise separable convolutional resnet with squeeze-and-excitation blocks for small-footprint keyword spotting." arXiv preprint arXiv:2004.12200 (2020). (Year: 2020). cited by examiner

XZhu, Chenzhuo, et al. "Trained ternary quantization." arXiv preprint arXiv:1612.01064 (2016). (Year: 2016). cited by examiner

Obukhov, Anton, et al. "T-basis: a compact representation for neural networks." International Conference on Machine Learning. PMLR, 2020. (Year: 2020). cited by examiner

Howard et al., "MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications," arXiv, Apr. 2017, 9 pages, retrieved from <https://www.semanticscholar.org/paper/MobileNets%3A-Efficient-Convolutional-Neural-Networks-Howard-Zhu/3647d6d0f151dc05626449ee09cc7bce55be497e>. cited by applicant

He et al., "Identity Mappings in Deep Residual Networks," ECCV, 2016, pp. 630-645. cited by applicant

Jaeger, H., "The "echo state" approach to analysing and training recurrent neural networks—with an Erratum note," German National Research Center for Information Technology GMD Technical Report, corrected version of the original 2001 paper, Jan. 26, 2010, 46 pages. cited by applicant

Kriman et al., "Quartznet: Deep automatic speech recognition with 1d time-channel separable convolutions," IEEE International Conference on Acoustics, Speech and Signal Processing, 2020, 5 pages, retrieved from <https://www.semanticscholar.org/paper/Quartznet%3A-Deep-Automatic-Speech-Recognition-with-Kriman-Beliaev/a0070675b4a7f55777cc1fceb5463c368e236337>. cited by applicant

Li et al., "Ternary weight networks," arXiv, 2016, 5 pages, retrieved from <https://arxiv.org/abs/1605.04711>. cited by applicant

Li et al., "Jasper: An End-to-End Convolutional Neural Acoustic Model," Interspeech, 2019, pp. 71-75. cited by applicant

Majumdar et al., "MatchboxNet: 1D Time-Channel Separable Convolutional Neural Network Architecture for Speech Commands Recognition," Interspeech, Oct. 2020, pp. 3356-3360. cited by applicant

Mordido et al., "Instant Quantization of Neural Networks using Monte Carlo Methods," EMC2: 5th Workshop on Energy Efficient Machine Learning and Cognitive Computing (NeurIPS 2019), 2019, 18 pages. cited by applicant

Zhu et al., "Trained Ternary Quantization," arXiv, 2016, 9 pages, retrieved from <https://arxiv.org/abs/1612.01064v1>. cited by applicant

He et al., "Deep Residual Learning for Image Recognition," IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2016, 12 pages, retrieved from <https://www.semanticscholar.org/paper/Deep-Residual-Learning-for-Image-Recognition-He-Zhang/2c03df8b48bf3fa39054345bafabfeff15bfd11d>. cited by applicant

Koluguri et al., "SpeakerNet: 1D Depth-wise Separable Convolutional Network for Text-Independent Speaker Recognition and Verification," arXiv: Audio and Speech Processing, 2020, 5 pages, retrieved from <https://www.semanticscholar.org/paper/SpeakerNet%3A-1D-Depth-wise-Separable-Convolutional-Koluguri-Li/888018da941ba3b4f38c1a28138f693b665fdeec>. cited by applicant

Warden et al., "Speech Commands: A Dataset for Limited-Vocabulary Speech Recognition," arXiv, 2018, 11 pages, retrieved <https://www.semanticscholar.org/paper/Speech-Commands%3A-A-Dataset-for-Limited-Vocabulary-Warden/da6e404d8911b0e5785019a79dc8607e0b313dc4>. cited by applicant

Panayotov et al., "Librispeech: An ASR Corpus Based on Public Domain Audio Books," IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), 2015, 5 pages, retrieved from <https://www.semanticscholar.org/paper/Librispeech%3A-An-ASR-corpus-based-on-public-domain-Panayotov-Chen/34038d9424ce602d7ac917a4e582d977725d4393>. cited by applicant

Graves et al., "Hybrid speech recognition with Deep Bidirectional LSTM," IEEE Workshop on Automatic Speech Recognition and Understanding, 2013, 6 pages, retrieved from <https://www.semanticscholar.org/paper/Hybrid-speech-recognition-with-Deep-Bidirectional-Graves-Jaitly/1149888d75af4ed5dfc25731b875651c3ccdeb2>. cited by applicant

Abdel-Hamid et al., "Convolutional Neural Networks for Speech Recognition," IEEE/ACM Transactions on Audio, Speech, and Language Processing, vol. 22, No. 10, Oct. 2014, pp. 1533-1545. cited by applicant

Sainath et al., "Convolutional, long short-term memory, fully connected deep neural networks," IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), 2015, 5 pages, retrieved from <https://www.semanticscholar.org/paper/Convolutional%2C-Long-Short-Term-Memory%2C-fully-Deep-Sainath-Vinyals/dec2ccce1ecb34bab02c42c2dd18cb468470adf8>. cited by applicant

Amodei et al., "Deep Speech 2 : End-to-End Speech Recognition in English and Mandarin," Proceedings of the 33rd International Conference on Machine Learning, 2016, 10 pages. cited by applicant

Bahdanau et al., "End-to-end Attention-Based Large Vocabulary Speech Recognition," IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), 2016, 8 pages, retrieved from <https://www.semanticscholar.org/paper/End-to-end-attention-based-large-vocabulary-speech-Bahdanau-Chorowski/878ba5458e9e51f0b341fd9117fa0b43ef4096d3>. cited by applicant

Dong et al., "Speech-Transformer: A No-Recurrence Sequence-to-Sequence Model for Speech Recognition," IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), 2018, pp. 5884-5888. cited by applicant

Bai et al., "An Empirical Evaluation of Generic Convolutional and Recurrent Networks for Sequence Modeling," arXiv, 2018, 14 pages, retrieved from <https://arxiv.org/abs/1803.01271>. cited by applicant

Sandler et al., "MobileNetV2: Inverted Residuals and Linear Bottlenecks," The IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2018, pp. 4510-4520, retrieved from <https://arxiv.org/abs/1801.04381v4>. cited by applicant

Tan et al., "EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks," International Conference on Machine Learning, 2019, 11 pages, retrieved from <https://arxiv.org/abs/1905.11946>. cited by applicant

Hayashi et al., "Einconv: Exploring Unexplored Tensor Network Decompositions for Convolutional Neural Networks," arXiv preprint, 2019, 12 pages. cited by applicant

Huang et al., "Cross-Language Transfer Learning, Continuous Learning, and Domain Adaptation for End-to-End Automatic Speech Recognition," arXiv preprint, 2020, 5 pages. cited by applicant

Molchanov et al., "Pruning Convolutional Neural Networks for Resource Efficient Transfer Learning," arXiv, 2016, 5 pages. cited by applicant

Povey et al., "Semi-Orthogonal Low-Rank Matrix Factorization for Deep Neural Networks," Interspeech, Sep. 2018, pp. 3743-3747. cited by applicant

Mordido et al., "Monte Carlo Gradient Quantization," CVPR 2020 Joint Workshop on Efficient Deep Learning in Computer Vision, 2020, 9 pages.

cited by applicant

Han et al., “Deep Compression: Compressing Deep Neural Network with Pruning, Trained Quantization and Huffman Coding,” arXiv preprint, 2015, 13 pages. cited by applicant

Huang et al., “Extreme learning machine: Theory and applications,” *Neurocomputing*, vol. 70, 2006, pp. 489-501. cited by applicant

Gallicchio et al., “Deep Randomized Neural Networks,” arXiv, 2020, 34 pages. cited by applicant

Huang et al., “Extreme Learning Machine for Regression and Multiclass Classification,” *IEEE Transactions on Systems, Man, and Cybernetics, Part B: Cybernetics*, vol. 42, No. 2, Apr. 2012, pp. 513-529. cited by applicant

Li et al., “On Feature Normalization and Data Augmentation,” arXiv, 2020, 13 pages. cited by applicant

Tang et al., “Acoustic Modeling with Densely Connected Residual Network for Multichannel Speech Recognition,” *Interspeech*, Sep. 2018, 7 pages. cited by applicant

Lin et al., “EdgeSpeechNets: Highly Efficient Deep Neural Networks for Speech Recognition on the Edge,” *32nd Conference on Neural Information Processing Systems (NIPS 2018)*, pp. 1-4. cited by applicant

Hannun et al., “Sequence-to-Sequence Speech Recognition with Time-Depth Separable Convolutions,” *Interspeech*, 2019, 5 pages, retrieved from <https://www.semanticscholar.org/paper/Sequence-to-Sequence-Speech-Recognition-with-Hannun-Lee/52fc149d40429bcf5090f956b343b779932215ed>. cited by applicant

Acero, A., “Acoustical and Environmental Robustness in Automatic Speech Recognition,” Thesis submitted to Department of Electrical and Computer Engineering Carnegie Mellon University, Sep. 13, 1990, 153 pages. cited by applicant

Skowronski et al., “Automatic speech recognition using a predictive echo state network classifier,” *Neural Networks*, vol. 20, 2007, pp. 414-423. cited by applicant

Boo et al., “Compression of Deep Neural Networks with Structured Sparse Ternary Coding,” *Journal of Signal Processing Systems*, vol. 91, 2019, pp. 1009-1019. cited by applicant

Krizhevsky et al., “ImageNet Classification with Deep Convolutional Neural Networks,” *Advances in Neural Information Processing Systems*, 2012, pp. 1-9, retrieved from https://proceedings.neurips.cc/paper_files/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf. cited by applicant

Kruithof et al., “Object recognition using deep convolutional neural networks with complete transfer and partial frozen layers,” *Proceeding of SPIE*, vol. 9995, 2016, 7 pages. cited by applicant

Gope et al., “Ternary MobileNets via Per-Layer Hybrid Filter Banks,” *IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 2020, pp. 3036-3046. cited by applicant

Kim et al., “X1000 Real-Time Phoneme Recognition VLSI Using Feed-Forward Deep Neural Networks,” *IEEE International Conference on Acoustic, Speech and Signal Processing (ICASSP)*, 2014, pp. 7510-7514. cited by applicant

Qian et al., “Do We Need Fully Connected Output Layers in Convolutional Networks?” arXiv, 2020, 8 pages, retrieved from <https://arxiv.org/pdf/2004.13587>. cited by applicant

Rongrong et al., “Deep Neural Network Compression and Acceleration: A Review,” *Journal of Computer Research and Development*, 2018, pp. 1871-1888. cited by applicant

Primary Examiner: Huang; Miranda M

Assistant Examiner: Black; Isis Marie

Attorney, Agent or Firm: Zilka-Kotab, P.C.

Background/Summary

CLAIM OF PRIORITY (1) This application claims the benefit of U.S. Provisional Application No. 63/106,287 titled “CONVOLUTIONS BY RANDOM TERNARY MATRICES,” filed Oct. 27, 2020, the entire contents of which is incorporated herein by reference.

FIELD OF THE INVENTION

(1) The present invention relates to artificial neural networks, and more particularly to incorporating one or more ternary matrices into a neural network.

BACKGROUND

(2) Artificial neural networks (ANNs) are commonly used computing systems that address a wide variety of tasks, such as classification, image recognition, regression, function approximation, samples of data according to a learned distribution, etc. However, ANN implementations often include fully connected layers (1×1 convolutions) that are computationally expensive and time-consuming to implement, train, and operate. There is a need for ANNs that are simpler, faster, and less resource-intensive to implement, train and operate.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 illustrates a flowchart of a method for incorporating a ternary matrix into a neural network, in accordance with an embodiment.
- (2) FIG. 2 illustrates a parallel processing unit, in accordance with an embodiment.
- (3) FIG. 3A illustrates a general processing cluster within the parallel processing unit of FIG. 2, in accordance with an embodiment.
- (4) FIG. 3B illustrates a memory partition unit of the parallel processing unit of FIG. 2, in accordance with an embodiment.
- (5) FIG. 4A illustrates the streaming multi-processor of FIG. 3A, in accordance with an embodiment.
- (6) FIG. 4B is a conceptual diagram of a processing system implemented using the PPU of FIG. 2, in accordance with an embodiment.
- (7) FIG. 4C illustrates an exemplary system in which the various architecture and/or functionality of the various previous embodiments may be implemented.
- (8) FIG. 5 is a conceptual diagram of a graphics processing pipeline implemented by the PPU of FIG. 2, in accordance with an embodiment.
- (9) FIG. 6 illustrates an exemplary directed graph of an artificial neural network, in accordance with an embodiment.
- (10) FIG. 7 illustrates an exemplary interpretation of projections onto halfspaces as ReLU neural units, in accordance with an embodiment.
- (11) FIG. 8 illustrates an exemplary comparison of two instances of dropout to an instance of partition, in accordance with an embodiment.
- (12) FIG. 9 illustrates exemplary selected paths in an artificial neural network, in accordance with an embodiment.
- (13) FIG. 10 illustrates subsampling activations of an entire feature layer online during inference (and training), in accordance with an embodiment.
- (14) FIG. 11 illustrates a flowchart of a method for generating paths to connect a set of neural units within an artificial neural network, in accordance with an embodiment.

(15) FIG. 12 illustrates a flowchart of a method for compressing an artificial neural network, in accordance with an embodiment.

(16) FIG. 13 illustrates a flowchart of a method for performing network normalization, in accordance with an embodiment.

DETAILED DESCRIPTION

(17) Artificial neural networks (ANNs) are computing systems that motivated by the human brain learn to perform tasks by considering examples. These ANNs are typically created by connecting several layers of artificial neurons using connections, where each artificial neuron may be connected to every other artificial neuron either directly or indirectly to create fully connected layers within the ANN. By substituting ternary matrices for one or more weight matrices of fully connected layers within the ANN, a complexity and resource usage of the ANN may be reduced, while improving the performance of the ANN.

(18) FIG. 1 illustrates a flowchart of a method **100** for incorporating a ternary matrix into a neural network, in accordance with an embodiment. Although method **100** is described in the context of a processing unit, the method **100** may also be performed by a program, custom circuitry, or by a combination of custom circuitry and a program. For example, the method **100** may be executed by a GPU (graphics processing unit), CPU (central processing unit), an FPGA (field programmable gate array), or any processor capable of performing ANN computations. Furthermore, persons of ordinary skill in the art will understand that any system that performs method **100** is within the scope and spirit of embodiments of the present invention.

(19) As shown in operation **102**, an artificial neural network (ANN) including a plurality of layers is created, where the plurality of layers includes at least one ternary matrix. In one embodiment, the ANN may include a plurality of paths each connecting at least one input of the ANN to at least one output of the ANN. In another embodiment, each of the plurality of paths may include a plurality of vertices each representing a neural unit within the ANN, and a plurality of edges each representing a weighted connection within the ANN.

(20) Additionally, in one embodiment, the ANN may include a convolutional model. For example, the ANN may transcribe speech, recognize speech commands, recognize speakers, etc. In another embodiment, the ANN may perform classification, image recognition, etc.

(21) In one embodiment, each ternary matrix may include a matrix that includes only the values -1 , 0 , and 1 . Further, in one embodiment each ternary matrix within the ANN may replace a fully connected layer within the ANN (e.g., a 1×1 convolution within the ANN, etc.).

(22) In another embodiment, each ternary matrix may be used instead of a 1×1 convolution. In yet another embodiment, each ternary matrix may be used instead of a 1×1 convolution in at least one residual layer of the ANN using at least one of a depth-wise separable convolution. In one embodiment, the depth-wise separable convolution is a time-channel separable convolution. In one example, output from a first fully-connected neural network layer may be provided as input to the ternary matrix, and output from the ternary matrix may be provided as input to a second fully-connected neural network layer.

(23) In another embodiment, each of the at least one ternary matrix may be randomly generated. For example, each of the ternary matrices may include a fixed sparse ternary random matrix. Further still, in one embodiment, the elements of at least one ternary matrix may be generated on-chip (e.g., utilizing one or more hardware processors that implement the ANN, etc.). In another embodiment, each ternary matrix may be generated by using a pseudo-random number generator. For example, the one or more hardware processors may initialize each of the ternary matrices at random. In another example, each of the values within the initialized matrices may be quantized utilizing a threshold value t to obtain ternary matrix values.

(24) For instance, all weights in the matrix may be quantized. In one embodiment, if a weight in the matrix is smaller than $-t$, then the weight is assigned a value of -1 . In another embodiment, if a weight in the matrix is between $-t$ and t , then the weight is assigned a value of 0 . In yet another embodiment, if a weight in the matrix is greater than t , then the weight is assigned a value of 1 . In another example, an absolute value of each of the values within the initialized matrices may be determined to obtain ternary matrix values.

(25) Also, in one embodiment, a hardware-based pseudo-random number generator may be used to generate the weights. In another embodiment, the threshold value t may be increased to increase the sparsity of the resulting ternary matrix. In yet another embodiment, the ternary matrix values may then be fixed within the ANN so that they do not change during the training of the ANN, during the performance of inference utilizing the ANN, etc.

(26) In addition, in one embodiment, for each ternary matrix, an initial state (e.g. one or more seed values, etc.) of the hardware-based pseudo-random number generator that was used to generate the weights for the quantization of the ternary matrix may be stored. This stored state may be used to regenerate/recreate the associated ternary matrix (e.g., in order to recreate an identical ANN at a later time, etc.).

(27) Furthermore, in one embodiment, each ternary matrix may generated by paths generated utilizing a low discrepancy sequence. In another embodiment, the plurality of layers may also include one or more identity matrices. For example, at least one skip link of the at least one residual layer of the ANN may be replaced by an identity matrix (e.g., a unit matrix, etc.). In another example, each identity matrix within the ANN may replace a fully connected layer within the ANN (e.g., a 1×1 convolution within the ANN, etc.). In another example, each identity matrix may copy the input directly to the output of the matrix (thereby eliminating the need for quadratic matrix multiplication during training/inference).

(28) Further still, as shown in operation **104**, the ANN is trained and inference is performed utilizing the trained ANN, where the at least one ternary matrix remains constant during the training and inference. For example, the values of each ternary matrix may remain constant and may not change during the training of the ANN. In one embodiment, the ANN may be trained utilizing labeled input training data. In another embodiment, the training may be supervised, semi-supervised, unsupervised, etc.

(29) Also, in one embodiment, the values of each ternary matrix may remain constant and may not change during both the training and the performance of inference utilizing the ANN. In another embodiment, the ANN may be used to perform inference by applying the trained ANN to an input data set to infer a result. For example, a data set may be input into the trained ANN, and the trained ANN may output one or more predictions based on the input.

(30) Additionally, in one embodiment, input data may be processed by the ANN to produce output data. For example, the input data may include one or more of image data, textual data, audio data, video data, random numbers, pseudo-random numbers, quasi-random numbers, low discrepancy sequences, etc. In another embodiment, the output data may include one or more of a classification, a categorization, a probability, a regression, function approximation, samples of data according to a learned distribution (e.g. generative adversarial networks (GANs)), etc.

(31) Further, in one embodiment, the input data may include environmental data (e.g., recorded image data of an environment surrounding an automobile, etc.), and the output data may include an identification/classification of one or more objects within the environmental data (such as cars, cyclists, pedestrians, etc.). In another embodiment, the ternary matrix may remain constant in that the values of the ternary matrix may not change (e.g., during the training of the ANN, during the performance of inference utilizing the ANN, etc.).

(32) Further still, in one embodiment, during the training of the ANN, and during the performance of inference utilizing the ANN, multiplication actions performed utilizing each of the at least one ternary matrix may be executed as a difference of two sums, where the first sum is the sum of all inputs to be weighted by the value 1 and the second sum is the sum of all inputs to be weighted by the value -1 .

(33) Also, in one embodiment, the ANN may be implemented as a hardware circuit. For example, during the performance of inference utilizing the ANN, multiplication actions performed utilizing each of the at least one ternary matrix may be executed as a difference of two sums. In another example, the difference may be amplified by an operational amplifier. In yet another example, each sum may be computed by wiring the selected input on either the positive or negative input line of the operational amplifier.

(34) In this way, a creation of the ANN may be simplified utilizing ternary and identity matrices, which are significantly less computationally expensive to create and execute when compared to fully connected layers/ 1×1 convolutions. This may reduce an amount of hardware resource usage needed to create/train/execute the ANN. Additionally, since ternary matrices may be created on-chip, only a seed value for the pseudo-random number generator may need to be stored to create such matrices, which may reduce an amount of hardware storage space needed to create the resulting ANN. Further, during an implementation of an ANN, each ternary matrix may be executed by calculating a difference of two sums, instead

of calculating more costly matrix multiplications of traditional fully-connected layers. This may improve a performance of ANNs utilizing ternary matrices.

(35) Further still, by increasing the threshold value t during quantization, the resulting ternary matrices may become more sparse, which may further improve the performance of ternary matrix execution (e.g., since more non-zero values exist in the resulting ternary matrix, and difference calculations need only be performed for non-zero values in the matrix).

(36) In yet another embodiment, the ANN may be created, trained, and/or implemented utilizing the parallel processing unit (PPU) **200** of FIG. 2.

(37) Additionally, since the number of nonzero weights may be selected in an embodiment, the computational complexity of a neural network layer may be reduced from quadratic to linear complexity. Neural networks may be represented by paths. In one embodiment, such path may be generated by sampling. Further, sampling may be performed proportional to discrete densities/weights within the ANN. Further still, sampling may be performed proportional to activations within the ANN. Also, weights of a neural network may be normalized, and the normalized weights may be propagated. In addition, network partitioning may be performed, and weights may be sub-sampled from a fully connected/convolutional neural network.

(38) More illustrative information will now be set forth regarding various optional architectures and features with which the foregoing framework may be implemented, per the desires of the user. It should be strongly noted that the following information is set forth for illustrative purposes and should not be construed as limiting in any manner. Any of the following features may be optionally incorporated with or without the exclusion of other features described. In one embodiment, the same random ternary matrices may be used throughout the network.

Parallel Processing Architecture

(39) FIG. 2 illustrates a parallel processing unit (PPU) **200**, in accordance with an embodiment. In an embodiment, the PPU **200** is a multi-threaded processor that is implemented on one or more integrated circuit devices. The PPU **200** is a latency hiding architecture designed to process many threads in parallel. A thread (i.e., a thread of execution) is an instantiation of a set of instructions configured to be executed by the PPU **200**. In an embodiment, the PPU **200** is a graphics processing unit (GPU) configured to implement a graphics rendering pipeline for processing three-dimensional (3D) graphics data in order to generate two-dimensional (2D) image data for display on a display device such as a liquid crystal display (LCD) device. In other embodiments, the PPU **200** may be utilized for performing general-purpose computations. While one exemplary parallel processor is provided herein for illustrative purposes, it should be strongly noted that such processor is set forth for illustrative purposes only, and that any processor may be employed to supplement and/or substitute for the same.

(40) One or more PPUs **200** may be configured to accelerate thousands of High Performance Computing (HPC), data center, and machine learning applications. The PPU **200** may be configured to accelerate numerous deep learning systems and applications including autonomous vehicle platforms, deep learning, high-accuracy speech, image, and text recognition systems, intelligent video analytics, molecular simulations, drug discovery, disease diagnosis, weather forecasting, big data analytics, astronomy, molecular dynamics simulation, financial modeling, robotics, factory automation, real-time language translation, online search optimizations, and personalized user recommendations, and the like.

(41) As shown in FIG. 2, the PPU **200** includes an Input/Output (I/O) unit **205**, a front end unit **215**, a scheduler unit **220**, a work distribution unit **225**, a hub **230**, a crossbar (Xbar) **270**, one or more general processing clusters (GPCs) **250**, and one or more partition units **280**. The PPU **200** may be connected to a host processor or other PPUs **200** via one or more high-speed NVLink **210** interconnect. The PPU **200** may be connected to a host processor or other peripheral devices via an interconnect **202**. The PPU **200** may also be connected to a local memory comprising a number of memory devices **204**. In an embodiment, the local memory may comprise a number of dynamic random access memory (DRAM) devices. The DRAM devices may be configured as a high-bandwidth memory (HBM) subsystem, with multiple DRAM dies stacked within each device.

(42) The NVLink **210** interconnect enables systems to scale and include one or more PPUs **200** combined with one or more CPUs, supports cache coherence between the PPUs **200** and CPUs, and CPU mastering. Data and/or commands may be transmitted by the NVLink **210** through the hub **230** to/from other units of the PPU **200** such as one or more copy engines, a video encoder, a video decoder, a power management unit, etc. (not explicitly shown). The NVLink **210** is described in more detail in conjunction with FIG. 4B.

(43) The I/O unit **205** is configured to transmit and receive communications (i.e., commands, data, etc.) from a host processor (not shown) over the interconnect **202**. The I/O unit **205** may communicate with the host processor directly via the interconnect **202** or through one or more intermediate devices such as a memory bridge. In an embodiment, the I/O unit **205** may communicate with one or more other processors, such as one or more the PPUs **200** via the interconnect **202**. In an embodiment, the I/O unit **205** implements a Peripheral Component Interconnect Express (PCIe) interface for communications over a PCIe bus and the interconnect **202** is a PCIe bus. In alternative embodiments, the I/O unit **205** may implement other types of well-known interfaces for communicating with external devices.

(44) The I/O unit **205** decodes packets received via the interconnect **202**. In an embodiment, the packets represent commands configured to cause the PPU **200** to perform various operations. The I/O unit **205** transmits the decoded commands to various other units of the PPU **200** as the commands may specify. For example, some commands may be transmitted to the front end unit **215**. Other commands may be transmitted to the hub **230** or other units of the PPU **200** such as one or more copy engines, a video encoder, a video decoder, a power management unit, etc. (not explicitly shown). In other words, the I/O unit **205** is configured to route communications between and among the various logical units of the PPU **200**.

(45) In an embodiment, a program executed by the host processor encodes a command stream in a buffer that provides workloads to the PPU **200** for processing. A workload may comprise several instructions and data to be processed by those instructions. The buffer is a region in a memory that is accessible (i.e., read/write) by both the host processor and the PPU **200**. For example, the I/O unit **205** may be configured to access the buffer in a system memory connected to the interconnect **202** via memory requests transmitted over the interconnect **202**. In an embodiment, the host processor writes the command stream to the buffer and then transmits a pointer to the start of the command stream to the PPU **200**. The front end unit **215** receives pointers to one or more command streams. The front end unit **215** manages the one or more streams, reading commands from the streams and forwarding commands to the various units of the PPU **200**.

(46) The front end unit **215** is coupled to a scheduler unit **220** that configures the various GPCs **250** to process tasks defined by the one or more streams. The scheduler unit **220** is configured to track state information related to the various tasks managed by the scheduler unit **220**. The state may indicate which GPC **250** a task is assigned to, whether the task is active or inactive, a priority level associated with the task, and so forth. The scheduler unit **220** manages the execution of a plurality of tasks on the one or more GPCs **250**.

(47) The scheduler unit **220** is coupled to a work distribution unit **225** that is configured to dispatch tasks for execution on the GPCs **250**. The work distribution unit **225** may track a number of scheduled tasks received from the scheduler unit **220**. In an embodiment, the work distribution unit **225** manages a pending task pool and an active task pool for each of the GPCs **250**. The pending task pool may comprise a number of slots (e.g., 32 slots) that contain tasks assigned to be processed by a particular GPC **250**. The active task pool may comprise a number of slots (e.g., 4 slots) for tasks that are actively being processed by the GPCs **250**. As a GPC **250** finishes the execution of a task, that task is evicted from the active task pool for the GPC **250** and one of the other tasks from the pending task pool is selected and scheduled for execution on the GPC **250**. If an active task has been idle on the GPC **250**, such as while waiting for a data dependency to be resolved, then the active task may be evicted from the GPC **250** and returned to the pending task pool while another task in the pending task pool is selected and scheduled for execution on the GPC **250**.

(48) The work distribution unit **225** communicates with the one or more GPCs **250** via XBar **270**. The XBar **270** is an interconnect network that couples many of the units of the PPU **200** to other units of the PPU **200**. For example, the XBar **270** may be configured to couple the work distribution unit **225** to a particular GPC **250**. Although not shown explicitly, one or more other units of the PPU **200** may also be connected to the XBar **270** via the hub **230**.

(49) The tasks are managed by the scheduler unit **220** and dispatched to a GPC **250** by the work distribution unit **225**. The GPC **250** is configured to

process the task and generate results. The results can be consumed by other tasks within the GPC 250, routed to a different GPC 250 via the XBar 270, or stored in the memory 204. The results can be written to the memory 204 via the partition units 280, which implement a memory interface for reading and writing data to/from the memory 204. The results can be transmitted to another PPU 200 or CPU via the NVLink 210. In an embodiment, the PPU 200 includes a number U of partition units 280 that is equal to the number of separate and distinct memory devices 204 coupled to the PPU 200. A partition unit 280 will be described in more detail below in conjunction with FIG. 3B.

(50) In an embodiment, a host processor executes a driver kernel that implements an application programming interface (API) that enables one or more applications executing on the host processor to schedule operations for execution on the PPU 200. In an embodiment, multiple compute applications are simultaneously executed by the PPU 200 and the PPU 200 provides isolation, quality of service (QoS), and independent address spaces for the multiple compute applications. An application may generate instructions (i.e., API calls) that cause the driver kernel to generate one or more tasks for execution by the PPU 200. The driver kernel outputs tasks to one or more streams being processed by the PPU 200. Each task may comprise one or more groups of related threads, referred to herein as a warp. In an embodiment, a warp comprises 32 related threads that may be executed in parallel. Cooperating threads may refer to a plurality of threads including instructions to perform the task and that may exchange data through shared memory. Threads and cooperating threads are described in more detail in conjunction with FIG. 4A.

(51) FIG. 3A illustrates a GPC 250 of the PPU 200 of FIG. 2, in accordance with an embodiment. As shown in FIG. 3A, each GPC 250 includes a number of hardware units for processing tasks. In an embodiment, each GPC 250 includes a pipeline manager 310, a pre-raster operations unit (PROP) 315, a raster engine 325, a work distribution crossbar (WDX) 380, a memory management unit (MMU) 390, and one or more Data Processing Clusters (DPCs) 320. It will be appreciated that the GPC 250 of FIG. 3A may include other hardware units in lieu of or in addition to the units shown in FIG. 3A.

(52) In an embodiment, the operation of the GPC 250 is controlled by the pipeline manager 310. The pipeline manager 310 manages the configuration of the one or more DPCs 320 for processing tasks allocated to the GPC 250. In an embodiment, the pipeline manager 310 may configure at least one of the one or more DPCs 320 to implement at least a portion of a graphics rendering pipeline. For example, a DPC 320 may be configured to execute a vertex shader program on the programmable streaming multiprocessor (SM) 340. The pipeline manager 310 may also be configured to route packets received from the work distribution unit 225 to the appropriate logical units within the GPC 250. For example, some packets may be routed to fixed function hardware units in the PROP 315 and/or raster engine 325 while other packets may be routed to the DPCs 320 for processing by the primitive engine 335 or the SM 340. In an embodiment, the pipeline manager 310 may configure at least one of the one or more DPCs 320 to implement a neural network model and/or a computing pipeline.

(53) The PROP unit 315 is configured to route data generated by the raster engine 325 and the DPCs 320 to a Raster Operations (ROP) unit, described in more detail in conjunction with FIG. 3B. The PROP unit 315 may also be configured to perform optimizations for color blending, organize pixel data, perform address translations, and the like.

(54) The raster engine 325 includes a number of fixed function hardware units configured to perform various raster operations. In an embodiment, the raster engine 325 includes a setup engine, a coarse raster engine, a culling engine, a clipping engine, a fine raster engine, and a tile coalescing engine. The setup engine receives transformed vertices and generates plane equations associated with the geometric primitive defined by the vertices. The plane equations are transmitted to the coarse raster engine to generate coverage information (e.g., an x,y coverage mask for a tile) for the primitive. The output of the coarse raster engine is transmitted to the culling engine where fragments associated with the primitive that fail a z-test are culled, and transmitted to a clipping engine where fragments lying outside a viewing frustum are clipped. Those fragments that survive clipping and culling may be passed to the fine raster engine to generate attributes for the pixel fragments based on the plane equations generated by the setup engine. The output of the raster engine 325 comprises fragments to be processed, for example, by a fragment shader implemented within a DPC 320.

(55) Each DPC 320 included in the GPC 250 includes an M-Pipe Controller (MPC) 330, a primitive engine 335, and one or more SMs 340. The MPC 330 controls the operation of the DPC 320, routing packets received from the pipeline manager 310 to the appropriate units in the DPC 320. For example, packets associated with a vertex may be routed to the primitive engine 335, which is configured to fetch vertex attributes associated with the vertex from the memory 204. In contrast, packets associated with a shader program may be transmitted to the SM 340.

(56) The SM 340 comprises a programmable streaming processor that is configured to process tasks represented by a number of threads. Each SM 340 is multi-threaded and configured to execute a plurality of threads (e.g., 32 threads) from a particular group of threads concurrently. In an embodiment, the SM 340 implements a SIMD (Single-Instruction, Multiple-Data) architecture where each thread in a group of threads (i.e., a warp) is configured to process a different set of data based on the same set of instructions. All threads in the group of threads execute the same instructions. In another embodiment, the SM 340 implements a SMT (Single-Instruction, Multiple Thread) architecture where each thread in a group of threads is configured to process a different set of data based on the same set of instructions, but where individual threads in the group of threads are allowed to diverge during execution. In an embodiment, a program counter, call stack, and execution state is maintained for each warp, enabling concurrency between warps and serial execution within warps when threads within the warp diverge. In another embodiment, a program counter, call stack, and execution state is maintained for each individual thread, enabling equal concurrency between all threads, within and between warps. When execution state is maintained for each individual thread, threads executing the same instructions may be converged and executed in parallel for maximum efficiency. The SM 340 will be described in more detail below in conjunction with FIG. 4A.

(57) The MMU 390 provides an interface between the GPC 250 and the partition unit 280. The MMU 390 may provide translation of virtual addresses into physical addresses, memory protection, and arbitration of memory requests. In an embodiment, the MMU 390 provides one or more translation lookaside buffers (TLBs) for performing translation of virtual addresses into physical addresses in the memory 204.

(58) FIG. 3B illustrates a memory partition unit 280 of the PPU 200 of FIG. 2, in accordance with an embodiment. As shown in FIG. 3B, the memory partition unit 280 includes a Raster Operations (ROP) unit 350, a level two (L2) cache 360, and a memory interface 370. The memory interface 370 is coupled to the memory 204. Memory interface 370 may implement 32, 64, 128, 1024-bit data buses, or the like, for high-speed data transfer. In an embodiment, the PPU 200 incorporates U memory interfaces 370, one memory interface 370 per pair of partition units 280, where each pair of partition units 280 is connected to a corresponding memory device 204. For example, PPU 200 may be connected to up to Y memory devices 204, such as high bandwidth memory stacks or graphics double-data-rate, version 5, synchronous dynamic random access memory, or other types of persistent storage.

(59) In an embodiment, the memory interface 370 implements an HBM2 memory interface and Y equals half U. In an embodiment, the HBM2 memory stacks are located on the same physical package as the PPU 200, providing substantial power and area savings compared with conventional GDDR5 SDRAM systems. In an embodiment, each HBM2 stack includes four memory dies and Y equals 4, with HBM2 stack including two 128-bit channels per die for a total of 8 channels and a data bus width of 1024 bits.

(60) In an embodiment, the memory 204 supports Single-Error Correcting Double-Error Detecting (SECCDED) Error Correction Code (ECC) to protect data. ECC provides higher reliability for compute applications that are sensitive to data corruption. Reliability is especially important in large-scale cluster computing environments where PPUs 200 process very large datasets and/or run applications for extended periods.

(61) In an embodiment, the PPU 200 implements a multi-level memory hierarchy. In an embodiment, the memory partition unit 280 supports a unified memory to provide a single unified virtual address space for CPU and PPU 200 memory, enabling data sharing between virtual memory systems. In an embodiment the frequency of accesses by a PPU 200 to memory located on other processors is traced to ensure that memory pages are moved to the physical memory of the PPU 200 that is accessing the pages more frequently. In an embodiment, the NVLink 210 supports address translation services allowing the PPU 200 to directly access a CPU's page tables and providing full access to CPU memory by the PPU 200.

(62) In an embodiment, copy engines transfer data between multiple PPUs 200 or between PPUs 200 and CPUs. The copy engines can generate page

faults for addresses that are not mapped into the page tables. In the memory partition unit **280** can then service the page faults, mapping the addresses into the page table, after which the copy engine can perform the transfer. In a conventional system, memory is pinned (i.e., non-pageable) for multiple copy engine operations between multiple processors, substantially reducing the available memory. With hardware page faulting, addresses can be passed to the copy engines without worrying if the memory pages are resident, and the copy process is transparent.

(63) Data from the memory **204** or other system memory may be fetched by the memory partition unit **280** and stored in the L2 cache **360**, which is located on-chip and is shared between the various GPCs **250**. As shown, each memory partition unit **280** includes a portion of the L2 cache **360** associated with a corresponding memory device **204**. Lower level caches may then be implemented in various units within the GPCs **250**. For example, each of the SMs **340** may implement a level one (L1) cache. The L1 cache is private memory that is dedicated to a particular SM **340**. Data from the L2 cache **360** may be fetched and stored in each of the L1 caches for processing in the functional units of the SMs **340**. The L2 cache **360** is coupled to the memory interface **370** and the XBar **270**.

(64) The ROP unit **350** performs graphics raster operations related to pixel color, such as color compression, pixel blending, and the like. The ROP unit **350** also implements depth testing in conjunction with the raster engine **325**, receiving a depth for a sample location associated with a pixel fragment from the culling engine of the raster engine **325**. The depth is tested against a corresponding depth in a depth buffer for a sample location associated with the fragment. If the fragment passes the depth test for the sample location, then the ROP unit **350** updates the depth buffer and transmits a result of the depth test to the raster engine **325**. It will be appreciated that the number of partition units **280** may be different than the number of GPCs **250** and, therefore, each ROP unit **350** may be coupled to each of the GPCs **250**. The ROP unit **350** tracks packets received from the different GPCs **250** and determines which GPC **250** that a result generated by the ROP unit **350** is routed to through the Xbar **270**. Although the ROP unit **350** is included within the memory partition unit **280** in FIG. 3B, in other embodiment, the ROP unit **350** may be outside of the memory partition unit **280**. For example, the ROP unit **350** may reside in the GPC **250** or another unit.

(65) FIG. 4A illustrates the streaming multi-processor **340** of FIG. 3A, in accordance with an embodiment. As shown in FIG. 4A, the SM **340** includes an instruction cache **405**, one or more scheduler units **410(K)**, a register file **420**, one or more processing cores **450**, one or more special function units (SFUs) **452**, one or more load/store units (LSUs) **454**, an interconnect network **480**, a shared memory/L1 cache **470**.

(66) As described above, the work distribution unit **225** dispatches tasks for execution on the GPCs **250** of the PPU **200**. The tasks are allocated to a particular DPC **320** within a GPC **250** and, if the task is associated with a shader program, the task may be allocated to an SM **340**. The scheduler unit **410(K)** receives the tasks from the work distribution unit **225** and manages instruction scheduling for one or more thread blocks assigned to the SM **340**. The scheduler unit **410(K)** schedules thread blocks for execution as warps of parallel threads, where each thread block is allocated at least one warp. In an embodiment, each warp executes 32 threads. The scheduler unit **410(K)** may manage a plurality of different thread blocks, allocating the warps to the different thread blocks and then dispatching instructions from the plurality of different cooperative groups to the various functional units (i.e., cores **450**, SFUs **452**, and LSUs **454**) during each clock cycle.

(67) Cooperative Groups is a programming model for organizing groups of communicating threads that allows developers to express the granularity at which threads are communicating, enabling the expression of richer, more efficient parallel decompositions. Cooperative launch APIs support synchronization amongst thread blocks for the execution of parallel algorithms. Conventional programming models provide a single, simple construct for synchronizing cooperating threads: a barrier across all threads of a thread block (i.e., the `syncthreads()` function). However, programmers would often like to define groups of threads at smaller than thread block granularities and synchronize within the defined groups to enable greater performance, design flexibility, and software reuse in the form of collective group-wide function interfaces.

(68) Cooperative Groups enables programmers to define groups of threads explicitly at sub-block (i.e., as small as a single thread) and multi-block granularities, and to perform collective operations such as synchronization on the threads in a cooperative group. The programming model supports clean composition across software boundaries, so that libraries and utility functions can synchronize safely within their local context without having to make assumptions about convergence. Cooperative Groups primitives enable new patterns of cooperative parallelism, including producer-consumer parallelism, opportunistic parallelism, and global synchronization across an entire grid of thread blocks.

(69) A dispatch unit **415** is configured to transmit instructions to one or more of the functional units. In the embodiment, the scheduler unit **410(K)** includes two dispatch units **415** that enable two different instructions from the same warp to be dispatched during each clock cycle. In alternative embodiments, each scheduler unit **410(K)** may include a single dispatch unit **415** or additional dispatch units **415**.

(70) Each SM **340** includes a register file **420** that provides a set of registers for the functional units of the SM **340**. In an embodiment, the register file **420** is divided between each of the functional units such that each functional unit is allocated a dedicated portion of the register file **420**. In another embodiment, the register file **420** is divided between the different warps being executed by the SM **340**. The register file **420** provides temporary storage for operands connected to the data paths of the functional units.

(71) Each SM **340** comprises L processing cores **450**. In an embodiment, the SM **340** includes a large number (e.g., 128, etc.) of distinct processing cores **450**. Each core **450** may include a fully-pipelined, single-precision, double-precision, and/or mixed precision processing unit that includes a floating point arithmetic logic unit and an integer arithmetic logic unit. In an embodiment, the floating point arithmetic logic units implement the IEEE 754-2008 standard for floating point arithmetic. In an embodiment, the cores **450** include 64 single-precision (32-bit) floating point cores, 64 integer cores, 32 double-precision (64-bit) floating point cores, and 8 tensor cores.

(72) Tensor cores configured to perform matrix operations, and, in an embodiment, one or more tensor cores are included in the cores **450**. In particular, the tensor cores are configured to perform deep learning matrix arithmetic, such as convolution operations for neural network training and inferencing. In an embodiment, each tensor core operates on a 4×4 matrix and performs a matrix multiply and accumulate operation $D=A \times B + C$, where A, B, C, and D are 4×4 matrices.

(73) In an embodiment, the matrix multiply inputs A and B are 16-bit floating point matrices, while the accumulation matrices C and D may be 16-bit floating point or 32-bit floating point matrices. Tensor Cores operate on 16-bit floating point input data with 32-bit floating point accumulation. The 16-bit floating point multiply requires 64 operations and results in a full precision product that is then accumulated using 32-bit floating point addition with the other intermediate products for a 4×4×4 matrix multiply. In practice, Tensor Cores are used to perform much larger two-dimensional or higher dimensional matrix operations, built up from these smaller elements. An API, such as CUDA 9 C++ API, exposes specialized matrix load, matrix multiply and accumulate, and matrix store operations to efficiently use Tensor Cores from a CUDA-C++ program. At the CUDA level, the warp-level interface assumes 16×16 size matrices spanning all 32 threads of the warp.

(74) Each SM **340** also comprises M SFUs **452** that perform special functions (e.g., attribute evaluation, reciprocal square root, and the like). In an embodiment, the SFUs **452** may include a tree traversal unit configured to traverse a hierarchical tree data structure. In an embodiment, the SFUs **452** may include texture unit configured to perform texture map filtering operations. In an embodiment, the texture units are configured to load texture maps (e.g., a 2D array of texels) from the memory **204** and sample the texture maps to produce sampled texture values for use in shader programs executed by the SM **340**. In an embodiment, the texture maps are stored in the shared memory/L1 cache **370**. The texture units implement texture operations such as filtering operations using mip-maps (i.e., texture maps of varying levels of detail). In an embodiment, each SM **240** includes two texture units.

(75) Each SM **340** also comprises N LSUs **454** that implement load and store operations between the shared memory/L1 cache **470** and the register file **420**. Each SM **340** includes an interconnect network **480** that connects each of the functional units to the register file **420** and the LSU **454** to the register file **420**, shared memory/L1 cache **470**. In an embodiment, the interconnect network **480** is a crossbar that can be configured to connect any of the functional units to any of the registers in the register file **420** and connect the LSUs **454** to the register file and memory locations in shared memory/L1 cache **470**.

(76) The shared memory/L1 cache **470** is an array of on-chip memory that allows for data storage and communication between the SM **340** and the primitive engine **335** and between threads in the SM **340**. In an embodiment, the shared memory/L1 cache **470** comprises 128 KB of storage capacity and is in the path from the SM **340** to the partition unit **280**. The shared memory/L1 cache **470** can be used to cache reads and writes. One or more of the shared memory/L1 cache **470**, L2 cache **360**, and memory **204** are backing stores.

(77) Combining data cache and shared memory functionality into a single memory block provides the best overall performance for both types of memory accesses. The capacity is usable as a cache by programs that do not use shared memory. For example, if shared memory is configured to use half of the capacity, texture and load/store operations can use the remaining capacity. Integration within the shared memory/L1 cache **470** enables the shared memory/L1 cache **470** to function as a high-throughput conduit for streaming data while simultaneously providing high-bandwidth and low-latency access to frequently reused data.

(78) When configured for general purpose parallel computation, a simpler configuration can be used compared with graphics processing. Specifically, the fixed function graphics processing units shown in FIG. 2, are bypassed, creating a much simpler programming model. In the general purpose parallel computation configuration, the work distribution unit **225** assigns and distributes blocks of threads directly to the DPCs **320**. The threads in a block execute the same program, using a unique thread ID in the calculation to ensure each thread generates unique results, using the SM **340** to execute the program and perform calculations, shared memory/L1 cache **470** to communicate between threads, and the LSU **454** to read and write global memory through the shared memory/L1 cache **470** and the memory partition unit **280**. When configured for general purpose parallel computation, the SM **340** can also write commands that the scheduler unit **220** can use to launch new work on the DPCs **320**.

(79) The PPU **200** may be included in a desktop computer, a laptop computer, a tablet computer, servers, supercomputers, a smart-phone (e.g., a wireless, hand-held device), personal digital assistant (PDA), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, and the like. In an embodiment, the PPU **200** is embodied on a single semiconductor substrate. In another embodiment, the PPU **200** is included in a system-on-a-chip (SoC) along with one or more other devices such as additional PPUs **200**, the memory **204**, a reduced instruction set computer (RISC) CPU, a memory management unit (MMU), a digital-to-analog converter (DAC), and the like.

(80) In an embodiment, the PPU **200** may be included on a graphics card that includes one or more memory devices **204**. The graphics card may be configured to interface with a PCIe slot on a motherboard of a desktop computer. In yet another embodiment, the PPU **200** may be an integrated graphics processing unit (iGPU) or parallel processor included in the chipset of the motherboard.

Exemplary Computing System

(81) Systems with multiple GPUs and CPUs are used in a variety of industries as developers expose and leverage more parallelism in applications such as artificial intelligence computing. High-performance GPU-accelerated systems with tens to many thousands of compute nodes are deployed in data centers, research facilities, and supercomputers to solve ever larger problems. As the number of processing devices within the high-performance systems increases, the communication and data transfer mechanisms need to scale to support the increased bandwidth.

(82) FIG. 4B is a conceptual diagram of a processing system **400** implemented using the PPU **200** of FIG. 2, in accordance with an embodiment. The exemplary system **465** may be configured to implement the method **100** shown in FIG. 1. The processing system **400** includes a CPU **430**, switch **410**, and multiple PPUs **200** each and respective memories **204**. The NVLink **210** provides high-speed communication links between each of the PPUs **200**. Although a particular number of NVLink **210** and interconnect **202** connections are illustrated in FIG. 4B, the number of connections to each PPU **200** and the CPU **430** may vary. The switch **410** interfaces between the interconnect **202** and the CPU **430**. The PPUs **200**, memories **204**, and NVLinks **210** may be situated on a single semiconductor platform to form a parallel processing module **425**. In an embodiment, the switch **410** supports two or more protocols to interface between various different connections and/or links.

(83) In another embodiment (not shown), the NVLink **210** provides one or more high-speed communication links between each of the PPUs **200** and the CPU **430** and the switch **410** interfaces between the interconnect **202** and each of the PPUs **200**. The PPUs **200**, memories **204**, and interconnect **202** may be situated on a single semiconductor platform to form a parallel processing module **425**. In yet another embodiment (not shown), the interconnect **202** provides one or more communication links between each of the PPUs **200** and the CPU **430** and the switch **410** interfaces between each of the PPUs **200** using the NVLink **210** to provide one or more high-speed communication links between the PPUs **200**. In another embodiment (not shown), the NVLink **210** provides one or more high-speed communication links between the PPUs **200** and the CPU **430** through the switch **410**. In yet another embodiment (not shown), the interconnect **202** provides one or more communication links between each of the PPUs **200** directly. One or more of the NVLink **210** high-speed communication links may be implemented as a physical NVLink interconnect or either an on-chip or on-die interconnect using the same protocol as the NVLink **210**.

(84) In the context of the present description, a single semiconductor platform may refer to a sole unitary semiconductor-based integrated circuit fabricated on a die or chip. It should be noted that the term single semiconductor platform may also refer to multi-chip modules with increased connectivity which simulate on-chip operation and make substantial improvements over utilizing a conventional bus implementation. Of course, the various circuits or devices may also be situated separately or in various combinations of semiconductor platforms per the desires of the user. Alternately, the parallel processing module **425** may be implemented as a circuit board substrate and each of the PPUs **200** and/or memories **204** may be packaged devices. In an embodiment, the CPU **430**, switch **410**, and the parallel processing module **425** are situated on a single semiconductor platform.

(85) In an embodiment, the signaling rate of each NVLink **210** is 20 to 25 Gigabits/second and each PPU **200** includes six NVLink **210** interfaces (as shown in FIG. 4B, five NVLink **210** interfaces are included for each PPU **200**). Each NVLink **210** provides a data transfer rate of 25 Gigabytes/second in each direction, with six links providing 300 Gigabytes/second. The NVLinks **210** can be used exclusively for PPU-to-PPU communication as shown in FIG. 4B, or some combination of PPU-to-PPU and PPU-to-CPU, when the CPU **430** also includes one or more NVLink **210** interfaces.

(86) In an embodiment, the NVLink **210** allows direct load/store/atomic access from the CPU **430** to each PPU's **200** memory **204**. In an embodiment, the NVLink **210** supports coherency operations, allowing data read from the memories **204** to be stored in the cache hierarchy of the CPU **430**, reducing cache access latency for the CPU **430**. In an embodiment, the NVLink **210** includes support for Address Translation Services (ATS), allowing the PPU **200** to directly access page tables within the CPU **430**. One or more of the NVLinks **210** may also be configured to operate in a low-power mode.

(87) FIG. 4C illustrates an exemplary system **465** in which the various architecture and/or functionality of the various previous embodiments may be implemented. The exemplary system **465** may be configured to implement the method **100** shown in FIG. 1.

(88) As shown, a system **465** is provided including at least one central processing unit **430** that is connected to a communication bus **475**. The communication bus **475** may be implemented using any suitable protocol, such as PCI (Peripheral Component Interconnect), PCI-Express, AGP (Accelerated Graphics Port), HyperTransport, or any other bus or point-to-point communication protocol(s). The system **465** also includes a main memory **440**. Control logic (software) and data are stored in the main memory **440** which may take the form of random access memory (RAM).

(89) The system **465** also includes input devices **460**, the parallel processing system **425**, and display devices **445**, i.e. a conventional CRT (cathode ray tube), LCD (liquid crystal display), LED (light emitting diode), plasma display or the like. User input may be received from the input devices **460**, e.g., keyboard, mouse, touchpad, microphone, and the like. Each of the foregoing modules and/or devices may even be situated on a single semiconductor platform to form the system **465**. Alternately, the various modules may also be situated separately or in various combinations of semiconductor platforms per the desires of the user.

(90) Further, the system **465** may be coupled to a network (e.g., a telecommunications network, local area network (LAN), wireless network, wide area network (WAN) such as the Internet, peer-to-peer network, cable network, or the like) through a network interface **435** for communication

purposes.

(91) The system **465** may also include a secondary storage (not shown). The secondary storage includes, for example, a hard disk drive and/or a removable storage drive, representing a floppy disk drive, a magnetic tape drive, a compact disk drive, digital versatile disk (DVD) drive, recording device, universal serial bus (USB) flash memory. The removable storage drive reads from and/or writes to a removable storage unit in a well-known manner.

(92) Computer programs, or computer control logic algorithms, may be stored in the main memory **440** and/or the secondary storage. Such computer programs, when executed, enable the system **465** to perform various functions. The memory **440**, the storage, and/or any other storage are possible examples of computer-readable media.

(93) The architecture and/or functionality of the various previous figures may be implemented in the context of a general computer system, a circuit board system, a game console system dedicated for entertainment purposes, an application-specific system, and/or any other desired system. For example, the system **465** may take the form of a desktop computer, a laptop computer, a tablet computer, servers, supercomputers, a smart-phone (e.g., a wireless, hand-held device), personal digital assistant (PDA), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, a mobile phone device, a television, workstation, game consoles, embedded system, and/or any other type of logic.

(94) While various embodiments have been described above, it should be understood that they have been presented by way of example only, and not limitation. Thus, the breadth and scope of a preferred embodiment should not be limited by any of the above-described exemplary embodiments, but should be defined only in accordance with the following claims and their equivalents.

Graphics Processing Pipeline

(95) In an embodiment, the PPU **200** comprises a graphics processing unit (GPU). The PPU **200** is configured to receive commands that specify shader programs for processing graphics data. Graphics data may be defined as a set of primitives such as points, lines, triangles, quads, triangle strips, and the like. Typically, a primitive includes data that specifies a number of vertices for the primitive (e.g., in a model-space coordinate system) as well as attributes associated with each vertex of the primitive. The PPU **200** can be configured to process the graphics primitives to generate a frame buffer (i.e., pixel data for each of the pixels of the display).

(96) An application writes model data for a scene (i.e., a collection of vertices and attributes) to a memory such as a system memory or memory **204**. The model data defines each of the objects that may be visible on a display. The application then makes an API call to the driver kernel that requests the model data to be rendered and displayed. The driver kernel reads the model data and writes commands to the one or more streams to perform operations to process the model data. The commands may reference different shader programs to be implemented on the SMs **340** of the PPU **200** including one or more of a vertex shader, hull shader, domain shader, geometry shader, and a pixel shader. For example, one or more of the SMs **340** may be configured to execute a vertex shader program that processes a number of vertices defined by the model data. In an embodiment, the different SMs **340** may be configured to execute different shader programs concurrently. For example, a first subset of SMs **340** may be configured to execute a vertex shader program while a second subset of SMs **340** may be configured to execute a pixel shader program. The first subset of SMs **340** processes vertex data to produce processed vertex data and writes the processed vertex data to the L2 cache **360** and/or the memory **204**. After the processed vertex data is rasterized (i.e., transformed from three-dimensional data into two-dimensional data in screen space) to produce fragment data, the second subset of SMs **340** executes a pixel shader to produce processed fragment data, which is then blended with other processed fragment data and written to the frame buffer in memory **204**. The vertex shader program and pixel shader program may execute concurrently, processing different data from the same scene in a pipelined fashion until all of the model data for the scene has been rendered to the frame buffer. Then, the contents of the frame buffer are transmitted to a display controller for display on a display device.

(97) FIG. 5 is a conceptual diagram of a graphics processing pipeline **500** implemented by the PPU **200** of FIG. 2, in accordance with an embodiment. The graphics processing pipeline **500** is an abstract flow diagram of the processing steps implemented to generate 2D computer-generated images from 3D geometry data. As is well-known, pipeline architectures may perform long latency operations more efficiently by splitting up the operation into a plurality of stages, where the output of each stage is coupled to the input of the next successive stage. Thus, the graphics processing pipeline **500** receives input data **501** that is transmitted from one stage to the next stage of the graphics processing pipeline **500** to generate output data **502**. In an embodiment, the graphics processing pipeline **500** may represent a graphics processing pipeline defined by the OpenGL® API. As an option, the graphics processing pipeline **500** may be implemented in the context of the functionality and architecture of the previous Figures and/or any subsequent Figure(s).

(98) As shown in FIG. 5, the graphics processing pipeline **500** comprises a pipeline architecture that includes a number of stages. The stages include, but are not limited to, a data assembly stage **510**, a vertex shading stage **520**, a primitive assembly stage **530**, a geometry shading stage **540**, a viewport scale, cull, and clip (VSCC) stage **550**, a rasterization stage **560**, a fragment shading stage **570**, and a raster operations stage **580**. In an embodiment, the input data **501** comprises commands that configure the processing units to implement the stages of the graphics processing pipeline **500** and geometric primitives (e.g., points, lines, triangles, quads, triangle strips or fans, etc.) to be processed by the stages. The output data **502** may comprise pixel data (i.e., color data) that is copied into a frame buffer or other type of surface data structure in a memory.

(99) The data assembly stage **510** receives the input data **501** that specifies vertex data for high-order surfaces, primitives, or the like. The data assembly stage **510** collects the vertex data in a temporary storage or queue, such as by receiving a command from the host processor that includes a pointer to a buffer in memory and reading the vertex data from the buffer. The vertex data is then transmitted to the vertex shading stage **520** for processing.

(100) The vertex shading stage **520** processes vertex data by performing a set of operations (i.e., a vertex shader or a program) once for each of the vertices. Vertices may be, e.g., specified as a 4-coordinate vector (i.e., $\langle x, y, z, w \rangle$) associated with one or more vertex attributes (e.g., color, texture coordinates, surface normal, etc.). The vertex shading stage **520** may manipulate individual vertex attributes such as position, color, texture coordinates, and the like. In other words, the vertex shading stage **520** performs operations on the vertex coordinates or other vertex attributes associated with a vertex. Such operations commonly including lighting operations (i.e., modifying color attributes for a vertex) and transformation operations (i.e., modifying the coordinate space for a vertex). For example, vertices may be specified using coordinates in an object-coordinate space, which are transformed by multiplying the coordinates by a matrix that translates the coordinates from the object-coordinate space into a world space or a normalized-device-coordinate (NCD) space. The vertex shading stage **520** generates transformed vertex data that is transmitted to the primitive assembly stage **530**.

(101) The primitive assembly stage **530** collects vertices output by the vertex shading stage **520** and groups the vertices into geometric primitives for processing by the geometry shading stage **540**. For example, the primitive assembly stage **530** may be configured to group every three consecutive vertices as a geometric primitive (i.e., a triangle) for transmission to the geometry shading stage **540**. In some embodiments, specific vertices may be reused for consecutive geometric primitives (e.g., two consecutive triangles in a triangle strip may share two vertices). The primitive assembly stage **530** transmits geometric primitives (i.e., a collection of associated vertices) to the geometry shading stage **540**.

(102) The geometry shading stage **540** processes geometric primitives by performing a set of operations (i.e., a geometry shader or program) on the geometric primitives. Tessellation operations may generate one or more geometric primitives from each geometric primitive. In other words, the geometry shading stage **540** may subdivide each geometric primitive into a finer mesh of two or more geometric primitives for processing by the rest of the graphics processing pipeline **500**. The geometry shading stage **540** transmits geometric primitives to the viewport SCC stage **550**.

(103) In an embodiment, the graphics processing pipeline **500** may operate within a streaming multiprocessor and the vertex shading stage **520**, the primitive assembly stage **530**, the geometry shading stage **540**, the fragment shading stage **570**, and/or hardware/software associated therewith, may sequentially perform processing operations. Once the sequential processing operations are complete, in an embodiment, the viewport SCC stage **550**

may utilize the data. In an embodiment, primitive data processed by one or more of the stages in the graphics processing pipeline **500** may be written to a cache (e.g., L1 cache, a vertex cache, etc.). In this case, in an embodiment, the viewport SCC stage **550** may access the data in the cache. In an embodiment, the viewport SCC stage **550** and the rasterization stage **560** are implemented as fixed function circuitry.

(104) The viewport SCC stage **550** performs viewport scaling, culling, and clipping of the geometric primitives. Each surface being rendered to is associated with an abstract camera position. The camera position represents a location of a viewer looking at the scene and defines a viewing frustum that encloses the objects of the scene. The viewing frustum may include a viewing plane, a rear plane, and four clipping planes. Any geometric primitive entirely outside of the viewing frustum may be culled (i.e., discarded) because the geometric primitive will not contribute to the final rendered scene. Any geometric primitive that is partially inside the viewing frustum and partially outside the viewing frustum may be clipped (i.e., transformed into a new geometric primitive that is enclosed within the viewing frustum. Furthermore, geometric primitives may each be scaled based on a depth of the viewing frustum. All potentially visible geometric primitives are then transmitted to the rasterization stage **560**.

(105) The rasterization stage **560** converts the 3D geometric primitives into 2D fragments (e.g., capable of being utilized for display, etc.). The rasterization stage **560** may be configured to utilize the vertices of the geometric primitives to setup a set of plane equations from which various attributes can be interpolated. The rasterization stage **560** may also compute a coverage mask for a plurality of pixels that indicates whether one or more sample locations for the pixel intersect the geometric primitive. In an embodiment, z-testing may also be performed to determine if the geometric primitive is occluded by other geometric primitives that have already been rasterized. The rasterization stage **560** generates fragment data (i.e., interpolated vertex attributes associated with a particular sample location for each covered pixel) that are transmitted to the fragment shading stage **570**.

(106) The fragment shading stage **570** processes fragment data by performing a set of operations (i.e., a fragment shader or a program) on each of the fragments. The fragment shading stage **570** may generate pixel data (i.e., color values) for the fragment such as by performing lighting operations or sampling texture maps using interpolated texture coordinates for the fragment. The fragment shading stage **570** generates pixel data that is transmitted to the raster operations stage **580**.

(107) The raster operations stage **580** may perform various operations on the pixel data such as performing alpha tests, stencil tests, and blending the pixel data with other pixel data corresponding to other fragments associated with the pixel. When the raster operations stage **580** has finished processing the pixel data (i.e., the output data **502**), the pixel data may be written to a render target such as a frame buffer, a color buffer, or the like.

(108) It will be appreciated that one or more additional stages may be included in the graphics processing pipeline **500** in addition to or in lieu of one or more of the stages described above. Various implementations of the abstract graphics processing pipeline may implement different stages. Furthermore, one or more of the stages described above may be excluded from the graphics processing pipeline in some embodiments (such as the geometry shading stage **540**). Other types of graphics processing pipelines are contemplated as being within the scope of the present disclosure. Furthermore, any of the stages of the graphics processing pipeline **500** may be implemented by one or more dedicated hardware units within a graphics processor such as PPU **200**. Other stages of the graphics processing pipeline **500** may be implemented by programmable hardware units such as the SM **340** of the PPU **200**.

(109) The graphics processing pipeline **500** may be implemented via an application executed by a host processor, such as a CPU. In an embodiment, a device driver may implement an application programming interface (API) that defines various functions that can be utilized by an application in order to generate graphical data for display. The device driver is a software program that includes a plurality of instructions that control the operation of the PPU **200**. The API provides an abstraction for a programmer that lets a programmer utilize specialized graphics hardware, such as the PPU **200**, to generate the graphical data without requiring the programmer to utilize the specific instruction set for the PPU **200**. The application may include an API call that is routed to the device driver for the PPU **200**. The device driver interprets the API call and performs various operations to respond to the API call. In some instances, the device driver may perform operations by executing instructions on the CPU. In other instances, the device driver may perform operations, at least in part, by launching operations on the PPU **200** utilizing an input/output interface between the CPU and the PPU **200**. In an embodiment, the device driver is configured to implement the graphics processing pipeline **500** utilizing the hardware of the PPU **200**.

(110) Various programs may be executed within the PPU **200** in order to implement the various stages of the graphics processing pipeline **500**. For example, the device driver may launch a kernel on the PPU **200** to perform the vertex shading stage **520** on one SM **340** (or multiple SMs **340**). The device driver (or the initial kernel executed by the PPU **300**) may also launch other kernels on the PPU **300** to perform other stages of the graphics processing pipeline **500**, such as the geometry shading stage **540** and the fragment shading stage **570**. In addition, some of the stages of the graphics processing pipeline **500** may be implemented on fixed unit hardware such as a rasterizer or a data assembler implemented within the PPU **300**. It will be appreciated that results from one kernel may be processed by one or more intervening fixed function hardware units before being processed by a subsequent kernel on an SM **340**.

Machine Learning

(111) Deep neural networks (DNNs) developed on processors, such as the PPU **200** have been used for diverse use cases, from self-driving cars to faster drug development, from automatic image captioning in online image databases to smart real-time language translation in video chat applications. Deep learning is a technique that models the neural learning process of the human brain, continually learning, continually getting smarter, and delivering more accurate results more quickly over time. A child is initially taught by an adult to correctly identify and classify various shapes, eventually being able to identify shapes without any coaching. Similarly, a deep learning or neural learning system needs to be trained in object recognition and classification for it get smarter and more efficient at identifying basic objects, occluded objects, etc., while also assigning context to objects.

(112) At the simplest level, neurons in the human brain look at various inputs that are received, importance levels are assigned to each of these inputs, and output is passed on to other neurons to act upon. An artificial neuron or perceptron is the most basic model of a neural network. In one example, a perceptron may receive one or more inputs that represent various features of an object that the perceptron is being trained to recognize and classify, and each of these features is assigned a certain weight based on the importance of that feature in defining the shape of an object.

(113) A deep neural network (DNN) model includes multiple layers of many connected perceptrons (e.g., nodes) that can be trained with enormous amounts of input data to quickly solve complex problems with high accuracy. In one example, a first layer of the DNN model breaks down an input image of an automobile into various sections and looks for basic patterns such as lines and angles. The second layer assembles the lines to look for higher level patterns such as wheels, windshields, and mirrors. The next layer identifies the type of vehicle, and the final few layers generate a label for the input image, identifying the model of a specific automobile brand.

(114) Once the DNN is trained, the DNN can be deployed and used to identify and classify objects or patterns in a process known as inference. Examples of inference (the process through which a DNN extracts useful information from a given input) include identifying handwritten numbers on checks deposited into ATM machines, identifying images of friends in photos, delivering movie recommendations to over fifty million users, identifying and classifying different types of automobiles, pedestrians, and road hazards in driverless cars, or translating human speech in real-time.

(115) During training, data flows through the DNN in a forward propagation phase until a prediction is produced that indicates a label corresponding to the input. If the neural network does not correctly label the input, then errors between the correct label and the predicted label are analyzed, and the weights are adjusted for each feature during a backward propagation phase until the DNN correctly labels the input and other inputs in a training dataset. Training complex neural networks requires massive amounts of parallel computing performance, including floating-point multiplications and additions that are supported by the PPU **200**. Inferencing is less compute-intensive than training, being a latency-sensitive process where a trained neural network is applied to new inputs it has not seen before to classify images, translate speech, and generally infer new information.

(116) Neural networks heavily rely on matrix math operations, and complex multi-layered networks require tremendous amounts of floating-point performance and bandwidth for both efficiency and speed. With thousands of processing cores, optimized for matrix math operations, and delivering tens to hundreds of TFLOPS of performance, the PPU **200** is a computing platform capable of delivering performance required for deep neural network-based artificial intelligence and machine learning applications.

Compressing 1D Time-Channel Separable Convolutions Using Sparse Random Ternary Matrices

(117) Speech and command recognition tasks tend to have strict low-latency requirements, which may be challenging to meet while using deep neural networks. Specifically, the employment of such networks on edge devices, which are bandwidth, energy, and area constrained, is often too costly in practical applications. Sparsity and quantization techniques are viable solutions, reducing the model size and computational cost while being suited for existing hardware.

(118) With this motivation in mind, sparsity and quantization may be combined to make compact speech and command recognition models even more compact. In particular, residual networks containing time-channel separable convolutions may be leveraged. Despite such compact designs, most weights are on the 1×1 -convolutions in the residual paths, which are equivalent to fully-connected layers applied to every time step. The memory and computational cost of these models may be decreased by replacing trained floating-point weights by constant, sparse random ternary matrices.

(119) This approach enables (i) faster training, since the weights need not be updated, (ii) faster on-chip inference, due to sparse, low-precision operations with no multiplications, and (iii) higher compression rates, because weights may not be stored but just computed on-the-fly resulting in (iv) reduced memory bandwidth.

(120) Random Ternary Matrices

(121) By analyzing the weights of the 1×1 -convolutions in the sequence of residual blocks, it may be observed that they do not expose any obvious visual structure. Specifically, after applying trained ternarized quantization, the resulting trained matrices maintain model performance even though they can hardly be distinguished from random ternary matrices.

(122) Neural networks with random weights show powerful capabilities. Therefore, random ternarized matrices may be used instead of 1×1 -convolutions in the residual blocks by keeping them constant while training the rest of the network. For each neuron, the ternarized weights simply compute a scaled difference of sums:

(123) $\text{ReLU}(\text{BN}(\sum_{i \in I^+} \text{Math. } x_i - \sum_{i \in I^-} \text{Math. } x_i))$,

(124) where the set I^+ contains all the indices of the connections with the weight +1 and I^- contains all the indices of the connections with the weight -1. The union $I = I^+ \cup I^-$ may be sparse, i.e. may not contain the indices of all neurons in the previous layer. This scaled sum is a far cheaper operation than the traditional floating-point matrix multiplication.

(125) Similar to echo state networks, a sufficient routing capacity is enabled by having trained layers before and after the constant random ternary layers. This matches observations that note that starting quantization too early in the network degrades performance. With this in mind, only ternarization is performed, and the weights in the 1×1 -convolutions are not updated inside the residual blocks (which possess the majority of the weights).

(126) Since the weights of the preceding temporal-convolution and the subsequent batch normalization layers are still updated, a scaling of the inputs and outputs may be learned, respectively. Implicitly, the tensor product of the aforementioned scaling vectors may define the weights of the ternary matrix, empowering this approach. The ternary matrix in combination with batch normalization is reminiscent of operational amplifiers, computing a difference of sums that is amplified before applying a non-linearity, e.g. ReLU. Moreover, the algorithm is related to extreme learning machines, where a more theoretical foundation can be derived.

(127) Exemplary Implementation Details

(128) Given a randomly initialized matrix, the random (and constant) ternary matrix may be generated by applying a thresholding procedure. A parameter t then serves as a threshold, where all absolute values less than t are quantized to zero and the remaining values are quantized to -1 and +1 according to their sign. If the matrices are initialized using uniformly distributed samples from $[-1; 1]$, $t \in [0, 1]$ will be the expected fraction of sparsity of the matrix: $t=0$ yields a dense, binary matrix, $t=1$ yields a zero matrix, and $t \in]0, 1[$ yields a ternary matrix.

(129) An exemplary implementation of this approach is shown below (given a uniformly distributed weight from $[-1; 1]$ and a threshold $t \in [0, 1]$):
compute signs pos=(weight.data>0.).int 2 () neg=(weight.data<0.).int 2 () # compute sparsity mask mask=(torch. abs(weight.data)>t).int 2 () # make ternary weight.data=(pos-neg)*mask # do not update weight during training weight.requires_grad=False

Implementation Variants

(130) The above implementations may be widely applicable and suited to existing hardware, allowing for a range of efficient implementations. Specifically, the random ternary matrices may be computed on-the-fly on the chip without accessing memory for weights. For example, a pseudo-random number generator may be used to determine the ternary values, where, given an initial state, a specific matrix may be deterministically generated at any time. Alternatively, each ternary value may be determined by a hash function of the matrix entry coordinates, where different matrices may be realized by different hash functions.

(131) Specific structured sparsity patterns supported by certain accelerators may be applied instead, taking full advantage of hardware support. Further, compute acceleration in the ternary layers may be leveraged by custom layer implementations.

Implementing a Linear Neural Network

(132) Artificial neural networks (ANNs) are computing systems that imitate a human brain by learning to perform tasks by considering examples. These ANNs are typically created by connecting several layers of neural units using connections, where each neural unit is connected to every other neural unit either directly or indirectly to create fully connected layers within the ANN. However, by representing an artificial neural network utilizing paths from an input of the ANN to an output of the ANN, a complexity of the ANN may be reduced, and the ANN may be trained and implemented in a much faster manner when compared to fully connected layers within the ANN.

(133) In one embodiment, a linear neural network may be implemented. For example, an artificial neural network (ANN) may be created that is represented by a plurality of paths each connecting at least one input of the ANN to at least one output of the ANN. Additionally, each of the plurality of paths includes a plurality of vertices each representing a neural unit within the ANN, and a plurality of edges each representing a weighted connection within the ANN.

(134) In one embodiment, input data is processed by the ANN to produce output data. In another embodiment, the input data may include one or more of image data, textual data, audio data, video data, random numbers, pseudo-random numbers, or quasi-random numbers, etc. In another embodiment, an embedding function may be used to map input data to a vector of floating point values to be processed by the ANN. In another embodiment, an embedding may be represented by an artificial neural network itself. In another embodiment, the output data may include one or more of a classification, a categorization, a probability, a regression, function approximation, samples of data according to a learned distribution (e.g. generative adversarial networks (GANs)), etc. In yet another embodiment, the input data may include environmental data (e.g., recorded image data of an environment surrounding an automobile, etc.), and the output data may include an identification/classification of one or more objects within the environmental data (such as cars, cyclists, pedestrians, etc.).

(135) In yet another embodiment, the ANN is trained, utilizing labeled input training data. For example, the training may be semi-supervised, unsupervised, etc. In yet another embodiment, the ANN is at the heart of a reinforcement learning machine to take actions and approximate a value. In this case the ANN is trained in a semi-supervised way by performing a simulation of a Markov chain. In yet another way, the ANN is used to

predict the next item of data and is trained to be supervised way by providing training data that includes the next data to be predicted.

(136) Additionally, in one embodiment, each of the plurality of paths may include a sequence of paths. In another embodiment, the sequence of paths may be partitioned into a plurality of contiguous blocks to create a series of artificial neural networks represented by the sequence of paths. In yet another embodiment, each of the plurality of paths may be generated by performing one or more of random sampling, pseudo-random sampling, and quasi-random sampling. In another embodiment, the pseudo-random sampling and quasi-random sampling may be performed in hardware. In yet another embodiment parts of the paths may be given and missing parts may be generated as described before. In yet another embodiment, the paths may be generated path by path or generation by generation, i.e. simultaneously one step at a time for all paths.

(137) Further, in one embodiment, the sampling may be performed on another ANN. In another embodiment, the sampling may be performed on fully connected layers of an ANN in order to determine the plurality of paths. In yet another embodiment, the plurality of paths may include one or more of a subset of all possible paths within the fully connected layers of the ANN or the convolutional layers of the ANN.

(138) For example, the other ANN may include a plurality of layers, where each layer includes a grouping of neural units (e.g., vertices). If the other ANN is fully connected, each neural unit (vertex) within a layer is connected (via an edge) to all neural units of a preceding layer as well as all neural units of a subsequent layer within the other ANN. These connections are called edges. In this way, all neural units of a fully connected ANN are either directly or indirectly connected to each other.

(139) Additionally, the paths may be generated by sampling a subset of all connections (edges) between the layers of neural units within the other ANN. For instance, between an input layer and a first layer within the other ANN, an edge may be sampled that connects a vertex within the input layer to a vertex of the first layer. An edge may then be sampled that connects the vertex of the first layer to a vertex of a second layer within the other ANN. This may be continued until a complete path is sampled that connects, via edges, the vertex within the input layer to a vertex of the output layer of the ANN, utilizing one vertex from each intermediate layer of the ANN. This may be performed for all vertices of the input layer and all vertices of the output layer and each vertex of the input layer of the other ANN is connected via a complete path of edges via the intermediate layers of the ANN to a unique vertex of the output layer.

(140) Further still, in one embodiment, the network may be created by uniformly sampling paths, utilizing an arbitrary network graph. In another embodiment, the arbitrary network graph may include an untrained ANN. In yet another embodiment, the plurality of paths may be selected from the untrained ANN, and may be used to create another ANN that represents the untrained ANN. In still another embodiment, this other ANN may then be trained.

(141) Also, in one embodiment, the weighted connections may be initialized deterministically. For example, each of the weighted connections within the ANN may be initialized with a constant value. In another embodiment, the weighted connections may be initialized with a value from a low discrepancy sequence.

(142) In addition, in one embodiment, the ANN may be created by sampling the plurality of paths proportional to a plurality of given weights of layers of the ANN. In another embodiment, the weights of a layer may be subsampled individually or in their entirety. In another embodiment, sampling the plurality of paths may include subsampling convolutional weights. In yet another embodiment, convolutional weights may be subsampled individually per filter or in their entirety. In still another embodiment, the plurality of paths may be selected such that the resulting network may be computed without performing any weight multiplications.

(143) Furthermore, in one embodiment, the ANN is created by sampling the plurality of paths proportional to a plurality of given activations of layers of the ANN. In another embodiment, the plurality of paths is selected proportional to the plurality of given activations during at least one of a training of the ANN and an inference performed utilizing the ANN. In yet another embodiment, the plurality of paths may be selected in a manner proportional to an error during back propagation. In still another embodiment, the plurality of paths may be selected such that the resulting network may be computed without performing any activation multiplications.

(144) Further still, in one embodiment, the plurality of paths may be selected without performing any multiplications by weights, or without performing any multiplications by activations. In another embodiment, the ANN may be created by sampling the plurality of paths proportional to both a plurality of given weights and a plurality of given activations of layers of the ANN. In yet another embodiment, the ANN may be a recurrent network.

(145) In yet another embodiment, the ANN may be created, trained, and/or implemented utilizing the parallel processing unit (PPU) 200 of FIG. 2.

(146) In this way, neural networks may be represented utilizing paths. Additionally, a complexity of a neural network may be reduced from quadratic to linear. Further, sampling may be performed proportionally to discrete densities/weights within the ANN. Further still, sampling may be performed proportional to activations within the ANN. Also, weights of a neural network may be normalized, and the normalized weights may be propagated. In addition, network partitioning may be performed, and weights may be sub-sampled from a fully connected/convolutional neural network.

Artificial Neural Networks with Linear Complexity

(147) Overview

(148) The average human brain has about 10^{11} nerve cells, where each of them may be connected to up to 10^4 others. We therefore investigate the question of whether there are algorithms for artificial neural networks that are linear in the number of neural units. Combining artificial neural networks and Monte Carlo and quasi-Monte Carlo methods allows one to derive general sampling-based algorithms for inference, training, compression, and quantization that are of linear complexity and have state of the art performance.

(149) Introduction

(150) Artificial neural networks are composed of neural units that are intended to imitate biological neurons. Such a neural unit computes a non-linear function of the weighted sum of the signals at its inputs. If every mentioned neural unit was connected to every other neural unit, the computation effort would be of quadratic complexity. However, it is known that not every neuron is connected to every other neuron in the brain. Even more so, the number of inputs to a neuron is bounded by a constant much smaller than the actual number of neurons.

(151) Artificial neural networks may contain so-called fully connected layers, where each neural unit of a layer is connected to every neural unit of the next layer. A large number of experiments have shown that large numbers of connections have little contribution to the end result and may be pruned without affecting the function of the artificial neural network.

(152) The weighted sum at the heart of a neural unit may be considered a scalar product or quadrature rule. Interpreting neural units as numerical integration algorithms, the principles of Monte Carlo and quasi-Monte Carlo methods may be applied to derive algorithms of linear complexity.

(153) Applying the principle of discrete density simulation allows for deriving algorithms of linear complexity for artificial neural networks. Representing artificial neural networks by paths, they may be trained much faster than conventional representations, and in addition may be trained in a sparse from scratch manner. As a consequence of sparsity, deterministic initialization becomes possible. Additionally, an algorithm may be derived that subsamples an artificial neural network such that the weights are only $\{-1, 0, 1\}$ and its complexity is linear in time and space in the number of neural units. Also, online sparsity may be obtained by subsampling the activations.

(154) Simulating Discrete Densities

(155) Scalar products are the defining operation of Hilbert spaces and are the basis of linear algorithms. They are at the core of machine learning and especially prominent in neural units, where a weighted sum of inputs is computed which is then passed as argument to a non-linear function (for example, see Equation 2 below).

(156) In what follows, we review how to evaluate scalar products using random and deterministic sampling, which is the basis for increasing the efficiency of artificial neural networks. See for example, "Massively parallel construction of radix tree forests for the efficient sampling of discrete probability distributions," by Binder et al., which discusses advanced algorithms for sampling discrete probability distributions, and which is herein

incorporated in its entirety.

(157) Without loss of generality, we assume a normalized set of weights, such that the absolute values of the weights add up to one, i.e.:

$$(158) \sum_{k=0}^{n-1} w_k = 1.$$

(159) If the weights are not normalized, we use weights

$$(160) \frac{w_k}{\sum_{k=0}^{n-1} w_k}$$

instead and assume that $\|w\|_1 \neq 0$.

(161) The absolute values of the weights then form a partition of the unit interval, where

(162)

$$P_m := \sum_{k=1}^m w_k \text{ and } 0 = P_0 < P_1 < \dots < P_{n-1} = 1: 0 = P_0 < w_1 < P_1 < w_2 < P_2 < \dots < P_n.$$

(163) Given N uniformly distributed samples $x_i \in [0, 1]$, this procedure allows one to approximate a scalar product

$$(164) \sum_{k=0}^{n-1} w_k a_k \approx \frac{1}{N} \sum_{i=0}^{N-1} \text{sign}(w_{j_i}) a_{j_i}, \quad (1)$$

where the index $j_i \in \{0, \dots, n-1\}$ is uniquely determined by $P_{j_i-1} \leq x_i < P_{j_i}$, as this satisfies that j_i is chosen with the probability $|w_{j_i}|$. $a = (a_0, \dots, a_{n-1})$ is a given vector of inputs.

(165) If the x_i are random numbers, then the procedure is the Monte Carlo integration of a scalar product by simulating the discrete density given by the absolute values of the normalized weights. It amounts to quantizing the weights to values in $\{-1, 0, +1\}$.

(166) The quality of the approximation may be improved by selecting samples x_i with improved uniform distribution. A deterministic choice would be low discrepancy sequences and a simple randomized alternative is jittered equidistant sampling, where:

$$(167) x_i = \frac{i}{N} + \xi_i$$

for one realization of the random variable $\xi_i \in [0, 1)$. Sorting the weights prior to taking their absolute values improves the approximation quality.

(168) Note that larger weights may be sampled multiple times. This case may be optimized by counting how many references were made and multiplying the corresponding activation by the number of references instead of looping over the references.

(169) On the other hand, duplicate references may rarely happen unless the weights vary largely. A crude approximation that ignores multiple references may work well in such cases. Alternatively, one may search a number of samples such that the resulting quantization is truly ternary (or binary) without duplicates. In another embodiment, the number of maximum bits to count duplicates may be fixed and samples may be added until this maximum representable count is reached. In yet another embodiment, counting may be continued beyond the maximum representable count by accepting a resulting truncation error.

(170) Sampling according to a distribution also may be implemented on parallel systems and if the weights are only available sequentially. Due to cancellation and rounding, an implementation in floating point numbers may not necessarily yield $\|w\|_1 = 1$ and care has to be taken to catch such cases. A practical workaround is to take the actual sum to scale the random variable.

(171) Artificial Neural Networks

(172) An example directed graph **600** of an artificial neural network is illustrated in FIG. 6. As shown, given the input vector $a_0 \dots a_{N-1}$, an output vector a_L is determined by the neural units $a_{l,i}$ **606A-N** and **608A-N**. Each neural unit in a layer l computes a weighted sum of the activations in the previous layer $l-1$ and applies an activation function that fixes its output.

(173) There are L layers with n_l neural units in l -th layer. Given weights $w_{l,k,i}$ and inputs activations $a_{0,k}$ the output of a neural unit is computed as:

$$(174) a_{l,i} := \max\{0, \sum_{k=0}^{n_{l-1}-1} w_{l,k,i} a_{l-1,k} - b_{l,i}\}. \quad (2)$$

(175) Thus one neural unit computes a weighted sum of the input activations, and applies a non-linear function to the weighted sum. The term $b_{l,i}$ is called bias and together with ReLU may be understood as a threshold compared to the weighted sum for activation.

(176) Properties of the Rectified Linear Unit (ReLU)

(177) While there are many choices for the non-linear output function of a neural unit, we focus on:

$\text{ReLU}(x) := \max\{0, x\}$,

and review some properties of neural units using the rectified linear unit function (ReLU). In fact, the ReLU activation function may be at the core of many functions used in artificial neural networks.

Leaky ReLU

(178) For example, the leaky ReLU is defined by:

$$\text{ReLU}(x) - \alpha \text{ReLU}(-x).$$

Max-Pooling

(179) For $\alpha = -1$ the leaky ReLU results to be the absolute value

$$|x| = \text{ReLU}(x) + \text{ReLU}(-x),$$

which in turn may be used to represent the maximum of two values:

$$(180) \max\{x, y\} = \frac{x+y}{2} + \frac{x-y}{2} \text{ReLU}(\frac{x+y}{2} + \frac{x-y}{2}) = \frac{1}{2} (x+y + \text{ReLU}(x-y) + \text{ReLU}(y-x))$$

(181) By induction, the maximum of two values may be extended to the maximum of an arbitrary number of values. More importantly, the identities show that computing the maximum may be represented by the ReLU activation function and thus may represent an optimization in an artificial neural network that does not need to be learned. A second consequence of this representation is the introduction of skip links, i.e. passing an input through one layer of a neural network.

(182) Residual Layers

(183) There is a strong connection between neural units and projection methods. Given a half space H^+ defined by a normal vector $\hat{w}$ and a perpendicular distance from the origin b , the projection of a point x onto that half space is given by

$$(184) P_H(x) := x - \min\{0, \hat{w} \cdot x - b\} \hat{w} = x - \text{ReLU}(\hat{w} \cdot x - b) \hat{w}. \quad (3)$$

(185) This is illustrated in FIG. 7, which is an interpretation **700** of projections onto halfspaces as ReLU neural units. As shown, the halfspace H^+ **702** is defined by the vector of weights $\hat{w}$ **704** as normal and a bias term b **706**, which represents the distance of the plane from the origin O **708**. Since $x_1 \in H^+$, the x_1 **710** is moved onto the boundary of H^+ **702** along the uniquely defined shortest distance vector, i.e. along the direction $\hat{w}$ **704**. For x_2 nothing is to be done, as already $x_2 \in H^+$ and hence $P_H(x_2) = x_2$.

(186) This may be used in adaptive projective subgradient methods for classification problems. In fact, the term $\text{ReLU}(\hat{w} \cdot x - b)$ translates into an if-clause that performs the projection, if the scalar product $\hat{w} \cdot x$ is larger than a threshold b .

(187) Note that usually the bias term is added (compare Equation 2). However, with the interpretation of the bias term as the distance from the origin, subtraction is more natural. The half space interpretation also provides some intuition about zero and non-zero bias: While non-zero bias allows one to construct finite convex sets, zero bias means that all planes of half spaces must intersect in the origin and hence the representable convex sets are infinite. Note that having one constant non-zero bias means that all planes have the same perpendicular distance from the origin.

(188) Strong relations of residual layers and ordinary differential equations are provided. Introducing a factor h , a residual layer may be written as:

$$(189) 0a_l = a_{l-1} + h \cdot \text{Math. } W_l^{(2)} \max\{0, W_l^{(1)} \cdot \text{Math. } a_{l-1}\} \Leftrightarrow \frac{a_l - a_{l-1}}{h} = W_l^{(2)} \max\{0, W_l^{(1)} \cdot \text{Math. } a_{l-1}\}$$

and after the equivalence transformation immediately looks like a step of the Euler method with step size h . With the transition $h \rightarrow \text{fwdarw.0}$, the left-hand side becomes $\{\text{dot over } (a)\}.\text{sub.l}$, which makes a residual layer an ordinary differential equation. This gave rise to trying different ordinary differential equations to determine $W.\text{sub.l}.\text{sup.}(1)$ and $W.\text{sub.l}.\text{sup.}(2)$ on the right-hand side and to perform inference and training using classic solvers for ordinary differential equations.

Network Normalization

(190) Given a positive factor $f \in \text{custom character}.\text{sup.}+$, a neural unit with the ReLU activation function has a scaling property:

$$(191) \max\{0, \text{Math. } w_{l,k,i} a_{l-1,k} - b_{l,i}\} = f \cdot \text{Math. } \max\{0, \frac{\text{Math. } w_{l,k,i} a_{l-1,k} - b_{l,i}}{f}\} \quad (4)$$

(192) Neural units with an activation function that has this scaling property may be normalized by selecting the linear factor

$$(193) f = \text{Math. } w_{l,i} \cdot \text{Math. } \frac{1}{\sum_{k=0}^{n_{l-1}-1} w_{l,k,i}} + \text{Math. } b_{l,i} \cdot \text{Math. } \frac{1}{\sum_{k=0}^{n_{l-1}-1} w_{l,k,i}} = \text{Math. } w_{l,k,i} \cdot \text{Math. } \frac{1}{\sum_{k=0}^{n_{l-1}-1} w_{l,k,i}} + \text{Math. } b_{l,i} \cdot \text{Math. } \frac{1}{\sum_{k=0}^{n_{l-1}-1} w_{l,k,i}}$$

such that the absolute values sum up to one. Using

$$(194) f = \text{Math. } w_{l,i} \cdot \text{Math. } \frac{1}{\sum_{k=0}^{n_{l-1}-1} w_{l,k,i}} = \text{Math. } w_{l,k,i} \cdot \text{Math. } \frac{1}{\sum_{k=0}^{n_{l-1}-1} w_{l,k,i}}$$

normalizes the absolute value of the weights such that they form a discrete probability distribution. A whole network may be normalized by starting to propagate the linear factors from the first layer through the network in a feed-forward fashion: The weights of the next layer are updated by multiplying with the respective linear factors from the previous layer. The last layer then needs to store the terminal linear factors for the output neural units. This network normalization transformation is deterministic. Considering an artificial neural network as an approximation operator, network normalization allows one to control the operator norm.

(195) As shown above, the leaky ReLU, absolute value, and max pooling have a positive linear factor property. Hence artificial neural networks built from them may be normalized as described. Note that for the example of the projection methods in Equation 3, we would use the Euclidean norm

$$(196) f = \|w_{l,i}\|_2 := \text{Math. } \sqrt{\sum_{k=0}^{n_{l-1}-1} |w_{l,k,i}|^2}$$

to obtain the normal vectors perpendicular to the plane they describe. The Euclidean norm also has been used to separate the length of a weight vector from its direction in order to accelerate training. It is a replacement for batch normalization.

Linear Algorithms for Artificial Neural Networks

(197) In a fully connected artificial neural network, the number of weights

$$(198) n_w = \sum_{l=1}^L n_{l-1} \cdot \text{Math. } n_l$$

is equivalent to the number of connections, where $n.\text{sub.l}$ is the number of neural units in layer l . However, in the brain, by far not every neuron is connected to all other neurons. One way to achieve a complexity linear in the number of neural units

$$(199) n = \sum_{l=1}^L n_l$$

is to bound the number of connections to a neural unit by a constant. In a first set of experiments, we sample connections of trained artificial neural networks proportional to the weight distribution of the neural units as explained above:

Fully Connected Layers

(200) A sparse artificial neural network may be created by selecting a fraction of the connections of a fully connected network by sampling according to the weights of each neural unit. In one example, using only 12% of the most important weights allows one to reach 97.52% of the accuracy of the full model. This simple approach already illustrates the power of subsampling neural networks. For this experiment, we used equidistant jittered sampling with one random offset per neural unit.

(201) Note that sampling neural units independently may lead to neural units or inputs not being connected. In addition, sampling a fraction of the weights is not yet a linear algorithm, although it already emphasizes the potential of the sampling approach.

(202) Convolutional Layers

(203) For a LeNet architecture on CIFAR-10, our model's best test accuracy is 69.12%. Applying subsampling to the convolution weights, we are able to get 88% of accuracy of the full model at only 50% sampled weights. The gains are smaller as compared to subsampling fully connected layers, as the number of weights is much smaller and in fact is already bounded by a constant, which is the filter kernel size. Therefore, subsampling individual filter kernels may only pay off for larger filter kernels.

(204) However, instead of sampling the weights of each filter one after another, sampling across all weights of all filters of a convolutional layer very much improves the situation. This way the selected connections much better represent the ensemble of all filters of a convolution layer, allowing for increased sparsity.

(205) Partition Instead of Dropout

(206) Dropout is used to regularize the training of artificial neural networks. Given a probability

$$(207) p,$$

a neural unit is not considered during training if for a uniformly distributed random number ξ we have

$$(208) p > \xi.$$

Such random number may be efficiently simulated by a linear feedback shift register generator, for example.

(209) This procedure, however, does not guarantee the coverage of all neural units. Assigning a neural unit to exactly one partition $p = \lfloor \xi \cdot \text{Math. } P \rfloor$ out of $P \in \text{custom character}$ partitions uses fewer pseudo-random number generator calls and guarantees that all neural units are considered. See, for example, FIG. 8, which compares two instances of dropout **802** and **804** to an instance of partition **806**.

(210) Except for the improved algorithmic efficiency, nothing is changed as shown in Table 1:

$$(211) \text{TABLE-US-00001 TABLE 1 LeNet on CIFAR-10} \quad \begin{array}{c} \text{DropoutAverageover} \\ p = 1/2 \text{ to } 1/9 \end{array} \quad \begin{array}{c} \text{Partitions Average of } P = 2 \text{ to } 9 \\ \text{Mean accuracy } 0.6062 \text{ } 0.6057 \end{array}$$

StdDev accuracy 0.0106 0.009

Representing Artificial Neural Networks by Paths

(212) With dropout or drop connect, it may happen that the output of a neural unit is not used, or that a neural unit may not be connected to any input. The same may be true for just sampling connections of a neural unit. However, constructing paths from the inputs to the outputs of an artificial neural network, there cannot be dangling neural units, as all neural units on a path use at least one input and propagate their output. Therefore, the complexity is bounded by the depth of the artificial neural network multiplied by the number of paths and hence is linear in both space and time.

(213) FIG. 9 illustrates exemplary selected paths **902-908** in an artificial neural network **900**, according to one exemplary embodiment. As shown, the computational complexity of the selected paths **902-908** in both time and space is bounded by the number of selected paths **902-908** times the number L of layers (e.g., the depth of the artificial neural network **900**).

(214) Table 2 illustrates Exemplary C++ declarations and initialization of the variables for artificial neural networks represented by paths.

LayerOffset[l] = $\sum.\text{sub.i} = 0.\text{sup.l} - 1.n.\text{sub.l}$ is the index of the first element of layer l in the activation and error arrays. The paths are either generated as

random walks connecting neural units selected by a pseudo-random number generator or by a deterministic sequence, where $P[i][l]$ is the l -th component of the i -th vector of a low discrepancy sequence.

(215) TABLE-US-00002 TABLE 2 const int InputSize = 784; const int Paths = 4 * InputSize; // example number of const float InitialWeight = 0.01f; // example network topology const int Layers = 6, NeuronsPerLayer[] = {InputSize, 256, 256 int Path[Layers][Paths]; // path length is number of Layers float Weight [Layers][Paths]; // one weight per edge in the path float *a; // activations float *error; // errors for back-propagation int LayerOffset[Layers]; // base index for layers in activa LayerOffset [0] = NeuronsPerLayer [0]; for (int l = 1; l < Layers; ++l) { for (int i = 0; i < Paths; ++i) Weight [l] [i] = InitialWeight; LayerOffset [l] = LayerOffset [l - 1] + NeuronsPerLayer[l]; } a = new float [LayerOffset [Layers - 1]]; error = new float [LayerOffset [Layers - 1]]; int Offset = 0; for (int l = 0; l < Layers; ++l) { for (int i = 0; i < Paths; ++i) #ifdef LOW_DISCREPANCY Path[l] [i] = Offset + (int) (P[i] [l] * NeuronsPerLayer #else Path[l] [i] = Offset + (int) (drand48() * NeuronsPerLayer #endif Offset += NeuronsPerLayer[l]; }

(216) Table 3 illustrates an exemplary linear complexity C++ implementation of the inference (feed-forward) step of an artificial neural network defined by paths using the ReLU activation function, according to one embodiment. The implementation shows the explicit execution of the ReLU activation function. Omitting that loop and uncommenting the if-statement implicitly realizes the ReLU activation function in a more hardware-amenable form, as the sign bit of the activation is sufficient to mask the accumulation operations.

(217) TABLE-US-00003 TABLE 3 inline void fastFeedForward(float *input) { for (int i = 0; i < LayerOffset[0]; ++i) a[i] = input[i]; // alternatively copy bias [i] values to a[i] for (int i = LayerOffset [0]; i < LayerOffset [Layers - 1]; ++ a[i] = 0.0f; for (int l = 1; l < Layers; ++l) { for (int i = 0; i < Paths; ++i) //if (a [Path [l - 1] [i]] > 0.0f) // implicit ReLU a [Path [l] [i]] += Weight [l] [i] * a [Path [l - 1] [i]]; // explicit ReLU for (int i = LayerOffset [l - 1]; i < LayerOffset[l]; ++i) a[i] = (a[i] > 0.0f ? a[i] : 0.0f); } }

(218) It still may happen that some neural units are never touched or that some connections are sampled multiple times. While this cannot be avoided using randomized methods to trace the paths, there exist number-theoretic constructions that allow one to create connection patterns with guarantees on coverage, fan-in, and fan-out.

(219) In that spirit, a hardware architecture for sparse artificial neural networks may be introduced. One stage of the pipeline architecture corresponds to one layer of an artificial neural network and takes care of feed-forward, back-propagation, and weight update in that layer. An artificial neural network is specified by the number of neural units per layer and both fan-in and fan-out for all neural units of a layer. The connections between layers are established using deterministic permutations, which allows one to avoid duplicate connections and to guarantee coverage of all neural units. The architecture has not been designed for convolutional neural networks. For sufficiently small fan-in and fan-out, the resulting artificial neural networks expose structural sparsity. Larger sparse neural networks may outperform smaller but dense networks of about the same complexity.

(220) Similarly, artificial neural networks based on expander graphs would guarantee at least one possible path from each input to each output. With respect to computing n output feature layers from m input feature layers, a subset of $D < m$ input feature layers may be selected for each of the n convolutions to reduce the complexity of convolutional layers. The implementation, however, is based on a random sampling method that does not guarantee the connection property.

(221) In one embodiment, fan-in and fan-out depend on the number of paths and the number $n_{\text{sub},l}$ of neural units per layer. If the number of neural units and the number of connections between layers are powers of two, the Sobol' sequence guarantees that the fan-in and fan-out per neural unit is an integer and constant per layer. Instead of selecting the same number of paths between layers, any power of two number of edges between layers generated using the Sobol' sequence guarantees the same. In addition, fan-in and fan-out may be adapted by trajectory splitting. See, for example, "Quasi-Monte Carlo image synthesis in a nutshell" by Alexander Keller, which is herein incorporated by reference in its entirety.

(222) Using a deterministic low discrepancy sequence, $P[l][i]$ may be understood as a function evaluation of the l -th component of the i -th vector of such a sequence, in other words, an address generator that for many low discrepancy sequences efficiently may be implemented in hardware. Especially for sequences like the Sobol' sequence, contiguous blocks of addresses generated by one component are permutations and hence guarantee for collision free routing when connecting layers in parallel.

(223) Progressive Sparse Training

(224) Instead of generating subsets of an artificial neural network by dropout or partitioning, a subgraph may be generated as the union of the neural units touched by a contiguous block of paths (see, for example, FIG. 9) from a sequence of paths. This procedure then is iterated for each batch of the training. Other than dropout, dropconnect, or partitioning, the subgraphs are guaranteed to be connected. This way, training complexity is guaranteed to be linear per iteration.

(225) Besides random walks on the graph of the network, low discrepancy sequences may be applied to sample the paths through the network in a progressive manner. This construction is based on so-called (t, s) -sequences in base b , which in fact are sequences of (t, m, s) -nets in base b .

(226) For example, the deterministic Sobol' low discrepancy sequence is a (t, s) -sequence in base $b=2$ that creates s dimensional points in the unit hypercube $[0, 1)^s$ that are uniformly distributed. These points may be very efficiently computed. A hardware realization is simple, because the arithmetic is entirely based on finite fields with a power of 2 elements. Now the components of the Sobol' sequence have the property that each contiguous block of a power of 2 elements forms a permutation, which guarantees that layers in a neural network linked by a quasi-random walk using the Sobol's sequence are linked by sets of permutations.

(227) Using the low discrepancy branch in the algorithm in Table 2, generating the deterministic quasi-random walks is as simple as evaluating the $P[l][i]$ in Table 2 using the s -dimensional points instead of a pseudo-random number generator. Given the layer l and the index i of the path, the corresponding component of the Sobol' sequence and its scrambled variants may be computed by a simple specific circuit of Boolean logic, which allows for using the sequence without explicitly storing its points.

(228) As these permutations are progressive, contiguous blocks of quasi-random walks generated by the Sobol' sequence may be used to subsample neural networks instead of using dropout.

(229) Note that low discrepancy sequences may be randomized and/or scrambled, which may preserve their permutation properties. Note that this is not necessarily guaranteed when using pseudo-random number generators instead of low discrepancy sequences.

(230) Similar to computer graphics, tracing paths through networks may be done by starting at the inputs, tracing back from the outputs, or starting at both ends and connecting path segments.

(231) Training Sparse from Scratch

(232) Instead of progressively subsampling a neural network, one set of paths may be used to fix a sparse architecture, see for example FIG. 9. Training sparse from scratch, the complexity of both inference and training is linear in the number of neural units touched by the paths.

(233) Uniform Initialization

(234) Before training, the weights of an artificial neural network need to be initialized.

(235) Usually, Gaussian distributed random variables or small scale uniform random variables are used, which results in unpredictable output during the initial training phase. This may be avoided by training a neural network to approximate uniform distribution before learning from data.

(236) For recurrent artificial neural networks, starting with a uniformly scaled diagonal unit matrix for the hidden-to-hidden weights may work.

(237) However, initializing the weight matrices with a small positive constant works nicely when we subsample the artificial neural network as shown above. While this may seem counterintuitive at first sight, observe that only a subset of the connections is selected. Hence the weights of the selected connections are initialized to a constant, for example, the inverse of the number of connections of a neural unit in order to be normalized, while the non-selected connection weights implicitly are set to zero. Thus the corresponding weight vector in a fully connected layer would consist of randomly or quasi-randomly chosen indices with zero and non-zero values. This configuration is sufficiently uniformly distributed to allow for

convergence training.

(238) Sparse by Quantization after Training

(239) Given a trained artificial neural network, it is simple to compress this network to use weights in $\{-1, 0, 1\}$ by just sampling proportional to the weights of each neural unit as explained above. Starting from the outputs, paths are sampled backwards proportional to the weights of each neural unit visited along the path. This corresponds to importance sampling paths in a weighted graph.

(240) The resulting network then only uses addition and subtraction and no more multiplication by weights. Note that connections may be sampled more than once and hence may be replaced by one connection with the weights summed up.

(241) In fact, a derivation of ternary weights is provided, and at the same time explains, why binary weights of only $\{-1, 1\}$ or $\{0, 1\}$ did not necessarily work: They could not be derived as discrete density approximation. One exemplary scheme may be considered a compiler that translates a trained artificial neural network into a ternary one. With respect to duplicate connections there is an interesting observation: In our scheme, connections may be sampled more than once if they are sufficiently important given the number of samples. Such duplicate connections cannot appear during weight quantization when not allowing for multiple references, however, they may appear as duplicate filter kernels.

(242) Sparse Inference

(243) While subsampling weights improves efficiency for large filters, it may not help much if the number of weights is small. Instead, the activations of an entire feature layer may be subsampled online during inference (and training) as illustrated in FIG. 10. More specifically, in sparse inference, neural units are selected proportional to their activations. This principle applies to any kind of layer, for example fully connected layers 1002 or convolutional layers 1004. This way, computations may be reduced dramatically, as only a fraction or a constant number of the connections are selected during feed-forward.

(244) This may be applied to any kind of layer in general, is especially useful for convolutional layers, and may help when the static quantization is not sufficiently efficient.

(245) In our implementation, we sample proportional to the activations of a layer online during inference and count how often an activation (i.e. neural unit) is selected. These counters then are divided by $\|a\|_{\text{sub}.1}$ and the result is fed forward. As compared to Equation 1, we now sample proportional to the distribution of the activations $a_{\text{sub}.k}$, which for the ReLU case are non-negative.

(246) If the activations are quantized to $\{0, 1\}$, we will only sum up weights without any multiplication. The subsampling may be considered another kind of regularization. In the limit, fixing the input and repeatedly subsampling activations as described resembles a network, where neural units spike proportional to their activations.

(247) In a similar way, errors may be propagated back through the network by sampling proportional to the error of a layer.

(248) The techniques introduced herein above may be applied to recurrent connections. While not only applicable to the quantization of recurrent weights, we may have sparse recurrent connections from scratch and sparse activations.

(249) Besides sampling proportional to the activations of a layer, we also may sample proportional to the product of activations and weights during inference. This corresponds most to what in Monte Carlo methods is called importance sampling. The neural unit under consideration then just adds up the -1 for sampled connections with negative weights and $+1$ for sampled connections with positive weights. The efficiency of such an approach depends on the number of entries in the discrete probability density. Obviously, it only pays off if the number is not quadratic.

CONCLUSION

(250) We showed that dropout partitions realize dropout at only a fraction of the pseudo-random numbers used before. Sampling proportional to the weights of a neural unit provides a derivation of ternary weights and yields a simple algorithm to create artificial neural networks that only use addition and subtraction instead of weight multiplication. As a consequence, it is straightforward to quantize neural networks to integer weights without retraining. Replacing drop connect and dropout by random- and quasi-random walks to generate paths through artificial neural networks guarantees no dangling neural units and linear complexity in space and time for both inference and training. Generating the paths before training results in artificial neural networks sparse from scratch. Constructing sparse artificial neural networks with linear complexity bears the potential of artificial neural networks without max-pooling, i.e. without resolution reduction, and therefore more precision.

(251) FIG. 11 illustrates a flowchart of a method 1100 for generating paths to connect a set of neural units within an artificial neural network, in accordance with an embodiment. Although method 1100 is described in the context of a processing unit, the method 1100 may also be performed by a program, custom circuitry, or by a combination of custom circuitry and a program. For example, the method 1100 may be executed by a GPU (graphics processing unit), CPU (central processing unit), an FPGA (field programmable gate array), or any processor or reconfigurable processor capable of performing the evaluation and/or training of ANNs. Furthermore, persons of ordinary skill in the art will understand that any system that performs method 1100 is within the scope and spirit of embodiments of the present invention.

(252) As shown in operation 1102, a number $L \in \text{custom character of layers}$ is selected for an artificial neural network.

(253) Additionally, as shown in operation 1104, for each layer $l \in \{0, \dots, L-1\}$, a number $n_{\text{sub}.i} \in \text{custom character of neural units}$ are selected, where $0 \leq j < n_{\text{sub}.l}$.

(254) Further, as shown in operation 1106, a number of paths n and an L -dimensional sequence $x_{\text{sub}.i} = (x_{\text{sub}.i,0}, \dots, x_{\text{sub}.i,L-1})$ uniformly distributed on $[0, 1]_{\text{sup}.L}$ are selected.

(255) Further still, as shown in operation 1108, for each $i \in \{0, \dots, n-1\}$ neural units $a_{\text{sub}.l}, \lfloor x_{\text{sub}.i,l} \cdot \text{Math}.n_{\text{sub}.l} \rfloor$ and $a_{\text{sub}.l}, \lfloor x_{\text{sub}.i,l+1} \cdot \text{Math}.n_{\text{sub}.l+1} \rfloor$ are connected to create the paths to connect the set of neural units within the artificial neural network.

(256) In this way, the paths generated to connect the set of neural units within the artificial neural network are random walks within the artificial neural network generated by a deterministic low discrepancy sequence.

(257) FIG. 12 illustrates a flowchart of a method 1200 for compressing an artificial neural network, in accordance with an embodiment. Although method 1200 is described in the context of a processing unit, the method 1200 may also be performed by a program, custom circuitry, or by a combination of custom circuitry and a program. For example, the method 1200 may be executed by a GPU (graphics processing unit), CPU (central processing unit), an FPGA (field programmable gate array), or any processor or reconfigurable processor capable of performing the evaluation and/or training of ANNs. Furthermore, persons of ordinary skill in the art will understand that any system that performs method 1200 is within the scope and spirit of embodiments of the present invention.

(258) As shown in operation 1202, an artificial neural network (ANN) is compressed by subsampling a plurality of edges each representing a weighted connection within the ANN. In one embodiment, the ANN may include layers that are fully connected. In another embodiment, the ANN may include layers that are convolutional. In yet another embodiment, each of the plurality of edges may be subsampled by performing one or more of random sampling, pseudo-random sampling, and quasi-random sampling.

(259) For example, an ANN may include a plurality of layers, where each layer includes a grouping of neural units. If the ANN is fully connected, each neural unit within a layer is connected (via an edge) to all neural units of a preceding layer as well as all neural units of a subsequent layer within the ANN. These connections are called edges. In this way, all neural units of a fully connected ANN are either directly or indirectly connected to each other.

(260) Additionally, the ANN may be compressed by sampling (e.g. selecting, etc.) a subset of all connections (edges) between the layers of neural units. In one embodiment, the edges may be sampled randomly. In another embodiment, the edges may be sampled in a pseudo-random manner. In yet another embodiment, the edges may be sampled in a quasi-random manner. Results of the sampling may include a compressed ANN that includes one or more of the sampled edges, the neural units associated with the sampled edges, and one or more paths created utilizing the one or more sampled edges.

(261) FIG. 13 illustrates a flowchart of a method 1300 for performing network normalization, in accordance with an embodiment. Although method 1300 is described in the context of a processing unit, the method 1300 may also be performed by a program, custom circuitry, or by a combination of custom circuitry and a program. For example, the method 1300 may be executed by a GPU (graphics processing unit), CPU (central processing unit), an FPGA (field programmable gate array), or any processor or reconfigurable processor capable of performing the evaluation and/or training of ANNs. Furthermore, persons of ordinary skill in the art will understand that any system that performs method 1300 is within the scope and spirit of embodiments of the present invention.

(262) As shown in operation 1302, network normalization is performed by normalizing a plurality of weights of a first layer of an artificial neural network (ANN) with a linear factor. Additionally, as shown in operation 1304, the normalized plurality of weights is propagated to successive layers of the ANN by multiplying weights of the successive layers by the linear factor. Further, a last layer of the ANN stores resulting weights to scale outputs of the ANN. In one embodiment, the ANN includes one or more of a ReLU activation function, a leaky ReLU activation function, a maxpool activation function, and an absolute value activation function.

(263) While various embodiments have been described above, it should be understood that they have been presented by way of example only, and not limitation. Thus, the breadth and scope of a preferred embodiment should not be limited by any of the above-described exemplary embodiments but should be defined only in accordance with the following claims and their equivalents.

(264) The disclosure may be described in the general context of computer code or machine-useable instructions, including computer-executable instructions such as program modules, being executed by a computer or other machine, such as a personal data assistant or other handheld device. Generally, program modules including routines, programs, objects, components, data structures, etc., refer to code that perform particular tasks or implement particular abstract data types. The disclosure may be practiced in a variety of system configurations, including hand-held devices, consumer electronics, general-purpose computers, more specialty computing devices, etc. The disclosure may also be practiced in distributed computing environments where tasks are performed by remote-processing devices that are linked through a communications network.

(265) As used herein, a recitation of “and/or” with respect to two or more elements should be interpreted to mean only one element, or a combination of elements. For example, “element A, element B, and/or element C” may include only element A, only element B, only element C, element A and element B, element A and element C, element B and element C, or elements A, B, and C. In addition, “at least one of element A or element B” may include at least one of element A, at least one of element B, or at least one of element A and at least one of element B. Further, “at least one of element A and element B” may include at least one of element A, at least one of element B, or at least one of element A and at least one of element B.

(266) The subject matter of the present disclosure is described with specificity herein to meet statutory requirements. However, the description itself is not intended to limit the scope of this disclosure. Rather, the inventors have contemplated that the claimed subject matter might also be embodied in other ways, to include different steps or combinations of steps similar to the ones described in this document, in conjunction with other present or future technologies. Moreover, although the terms “step” and/or “block” may be used herein to connote different elements of methods employed, the terms should not be interpreted as implying any particular order among or between various steps herein disclosed unless and except when the order of individual steps is explicitly described.

Claims

1. A method comprising: at one or more devices: creating, by the one or more devices, an artificial neural network (ANN) including a plurality of layers, the plurality of layers including at least one ternary matrix; training the ANN, where the at least one ternary matrix remains constant during the training of the ANN; and performing an inference utilizing the ANN during which multiplication actions performed utilizing each of the at least one ternary matrix are executed as a difference of two sums, where a first sum is a sum of all inputs to be weighted by the value 1 and a second sum is the sum of all inputs to be weighted by the value -1 ; and outputting a result of the inference.
2. The method of claim 1, wherein each of the at least one ternary matrix includes only the values -1 , 0, and 1.
3. The method of claim 1, wherein the at least one ternary matrix replaces one or more fully connected layers within the ANN.
4. The method of claim 1, wherein each of the at least one ternary matrix is used instead of a 1×1 convolution in at least one layer within the ANN.
5. The method of claim 1, wherein each of the at least one ternary matrix is used instead of a 1×1 convolution in at least one residual layer of the ANN using at least one of a depth-wise separable convolution.
6. The method of claim 5, wherein the at least one depth-wise separable convolution includes a time-channel separable convolution.
7. The method of claim 5, wherein at least one skip link of the at least one residual layer of the ANN is replaced by a unit matrix.
8. The method of claim 1, wherein all elements of the at least one ternary matrix are generated on-chip.
9. The method of claim 1, wherein each of the at least one ternary matrix includes paths generated utilizing a low discrepancy sequence.
10. The method of claim 1, wherein each of the at least one ternary matrix is generated using a pseudo-random number generator.
11. The method of claim 1, further comprising, for each of the at least one ternary matrix: storing an initial state of a hardware-based pseudo-random number generator that was used to generate weights for a quantization to compute the ternary matrix, and using the stored state to regenerate the ternary matrix.
12. The method of claim 1, wherein the plurality of layers also include one or more identity matrices.
13. The method of claim 12, wherein each of the one or more identity matrices within the ANN replace a fully connected layer within the ANN.
14. A non-transitory computer-readable media storing computer instructions which when executed by one or more processors of a device cause the one or more processors to cause the device to: create an artificial neural network (ANN) including a plurality of layers, the plurality of layers including at least one ternary matrix; train the ANN, where the at least one ternary matrix remains constant during the training of the ANN; and perform an inference utilizing the ANN during which multiplication actions performed utilizing each of the at least one ternary matrix are executed as a difference of two sums, where a first sum is a sum of all inputs to be weighted by the value 1 and a second sum is the sum of all inputs to be weighted by the value -1 ; and output a result of the inference.
15. A computer system comprising: a hardware circuit implementing an artificial neural network (ANN), wherein the hardware circuit includes a plurality of layers, the plurality of layers including at least one ternary matrix; and a hardware processor that executes to train the ANN, where the at least one ternary matrix remains constant during the training of the ANN.
16. The computer system of claim 15, wherein: during a performance of inference utilizing the ANN, multiplication actions performed utilizing each of the at least one ternary matrix are executed as a difference of two sums, the difference is amplified by an operational amplifier, and each sum is computed by wiring a selected input on either a positive or negative input line of the operational amplifier.
17. A hardware circuit implementing an artificial neural network (ANN) comprising: a plurality of neural network layers including at least one first fully connected layer; a ternary matrix; and a unit matrix; wherein the ternary matrix is representative of a second fully-connected layer or is representative of a 1×1 convolution layer; wherein the unit matrix is representative of a skip link.
18. A hardware circuit implementing a neural network comprising: a first fully-connected neural network layer; a ternary matrix coupled to the first fully-connected neural network layer, wherein the ternary matrix is used instead of a 1×1 convolution in a residual layer of the neural network using at least one of a depth-wise separable convolution, and wherein at least one skip link of the residual layer of the neural network is replaced by a unit matrix; and a second fully-connected neural network layer coupled to at least one of the first fully-connected neural network layer and the ternary matrix.

19. The hardware circuit of claim 18, wherein output from the first fully-connected neural network layer is provided as input to the ternary matrix.
 20. The hardware circuit of claim 18, wherein output from the ternary matrix is provided as input to the second fully-connected neural network layer.
-