

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12400137
Kind Code	B1
Date of Patent	August 26, 2025
Inventor(s)	Vivekraja; Vignesh et al.

Bidirectional network on a data-flow centric processor

Abstract

Techniques for implementing bidirectional networks, such as bidirectional recurrent neural networks (RNNs), on data-flow centric processors are disclosed. An input to a bidirectional network is provided. The bidirectional network includes a forward network and a reverse network. The forward network performs a set of operations in a forward direction and the reverse network performs the set of operations in a reverse direction. A forward output is generated using the forward network based on the input. A first transformation is performed on the input to generate a transformed input. A reverse output is generated using the reverse network based on the transformed input. A second transformation is performed on the reverse output to generate a transformed output. The forward output and the transformed output are combined to produce an overall output.

Inventors: Vivekraja; Vignesh (Santa Clara, CA), Huang; Randy Renfu (Morgan Hill, CA), Gai; Jiading (Seattle, WA), Hah; Thiam Khean (Milpitas, CA)

Applicant: Amazon Technologies, Inc. (Seattle, WA)

Family ID: 1000004393528

Assignee: Amazon Technologies, Inc. (Seattle, WA)

Appl. No.: 16/588931

Filed: September 30, 2019

Publication Classification

Int. Cl.: G06N5/046 (20230101); G06N3/08 (20230101); G06N20/00 (20190101)

U.S. Cl.:

CPC G06N5/046 (20130101); G06N3/08 (20130101); G06N20/00 (20190101);

Field of Classification Search

CPC: G06N (5/046); G06N (3/08); G06N (20/00)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
10210860	12/2018	Ward	N/A	G10L 25/18
10733506	12/2019	Ogale	N/A	G06N 3/084
2016/0292567	12/2015	George	N/A	G06N 3/047
2017/0061328	12/2016	Majumdar	N/A	G06K 9/6232
2019/0206020	12/2018	Ould-Ahmed-Vall	N/A	G06F 9/3863
2020/0210839	12/2019	Lo	N/A	G06F 9/5027
2020/0293770	12/2019	Chen	N/A	G06N 3/045
2020/0294287	12/2019	Schlemper	N/A	G06V 10/431

OTHER PUBLICATIONS

Khorasani et al. "In-Register Parameter Caching for Dynamic Neural Nets with Virtual Persistent Processor Specialization", Dec. 13, 2018 51st Annual IEEE (Year: 2018). cited by examiner

Zhang et al. "Relation classification via recurrent neural network with attention and tensor layers", vol. 1, No. 3, pp. 234-244, Sep. 2018 (Year: 2018). cited by examiner

Kitza et al. ("Comparison of BLSTM-Layer-Specific Affine Transformations for Speaker Adaptation", Interspeech 2018) (Year: 2018). cited by examiner

Yokota et al ("Missing Slice Recovery for Tensors Using a Low-rank Model in Embedded Space", IEEE, 2018, pp. 8251-8259) (Year : 2018). cited by examiner

Chien et al. ("NPL, Tensor-Factorized Neural Networks", IEEE, vol. 29, No. 5, May 2018) (Year: 2018). cited by examiner

Zhang et al. ("Relation Classification via Recurrent Neural Network with Attention and Tensor Layers", pp. 234-244 vol. 1, No. 3, Sep. 2018) (Year: 2018). cited by examiner

Primary Examiner: Huntley; Michael J

Assistant Examiner: Kassim; Imad

Attorney, Agent or Firm: Kilpatrick Townsend & Stockton LLP

Background/Summary

BACKGROUND

(1) Artificial neural networks are computing systems with an architecture based on biological neural networks. Artificial neural networks can be trained using training data to learn how to perform a certain task, such as identifying or classifying physical objects, activities, characters, etc., from images or videos. An artificial neural network, such as a deep neural network, may include multiple layers of processing nodes. Each processing node in a layer can perform computations on input data generated by processing nodes in the preceding layer to generate output data. For example, a processing node may perform a set of arithmetic operations such as

multiplications and additions to generate an intermediate output, or perform post-processing operations on the intermediate output to generate a final output. An artificial neural network may include thousands or more of processing nodes and millions or more of parameters.

(2) The architecture of a neural network may include an input layer, an output layer, and a number of intermediate layers, often referred to as hidden layers. Each layer executes a computation on the outputs of the previous layer, with the last layer (the output layer) providing a final result. With more layers, a neural network can, theoretically, perform more complex tasks, such as language translations and identifying (or classifying) the contents of an image. A neural network with more than three hidden layers is sometimes referred to as a deep neural network. Deep neural networks can have many hidden layers, such as, for example, between five and more than a thousand layers.

(3) Neural networks can be implemented using a central processing unit (CPU) to perform the computations. CPUs, however, tend to be optimized for sequential rather than parallel computations, and thus can suffer from poor response times. Graphics processing units (GPUs) are optimized for parallel computations, but not necessarily for the result from one computation unit to be provided directly to another computation unit. Often, the result must first be written to a memory and then read back. Although GPUs can have better response times than CPUs, it would still be desirable to improve the execution time of a neural network. Recently, special-purpose processors, such as neural network accelerators, have been developed to execute neural networks more efficiently than either CPUs or GPUs.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) Various embodiments in accordance with the present disclosure will be described with reference to the drawings, in which:
- (2) FIG. 1 illustrates an example of a computational flow model for a neural network;
- (3) FIG. 2 illustrates an example of a unidirectional recurrent neural network (RNN) layer;
- (4) FIG. 3 illustrates an example of a bidirectional RNN layer;
- (5) FIGS. 4A and 4B illustrate tables that show the input states, input values, output states and output values for a bidirectional RNN layer;
- (6) FIGS. 5A and 5B illustrate tables that show the input states, input values, output states and output values for a bidirectional RNN layer;
- (7) FIG. 6 illustrates an example of a bidirectional RNN layer;
- (8) FIGS. 7A and 7B illustrate tables that show the input states, input values, output states and output values for a bidirectional RNN layer;
- (9) FIG. 8 illustrates a table that shows the input states, input values, output states and output values for a bidirectional RNN layer;
- (10) FIG. 9 illustrates an example of a possible first transformation performed by a first transformation layer;
- (11) FIG. 10 illustrates an example of a possible second transformation performed by a second transformation layer;
- (12) FIG. 11 includes a flowchart illustrating an example of a method for executing a bidirectional network on a data-flow centric processor;
- (13) FIG. 12 includes a block diagram illustrating an example of a host system **1200** on which a compiler, a data-flow centric processor, and a host processor, such as are described herein, can run;
- (14) FIG. 13 includes a block diagram that illustrates an example of an acceleration engine;
- (15) FIG. 14 is a block diagram illustrating an example of an integrated circuit device; and
- (16) FIG. 15 includes a diagram of an example network.

DETAILED DESCRIPTION

(17) For some types of neural networks, such as recurrent neural networks (RNNs), the operations performed by a particular layer or node are repeated multiple times, each time with different input data and/or different input states. During each iteration, the input data and/or the input state of a layer or node may be altered based on the output data and/or output state from the previous iteration. Such neural networks may be represented with a data flow graph having a feedback loop that indicates that the output data and/or output state from a particular layer/node is fed back as the input data and/or input state for the particular layer/node. The number of iterations that are performed on the set of operations may be static or dynamic, as dictated by the control flow of the neural network. For example, the number of iterations may be static if the set of operations are implemented in a “for” loop with no break conditions. In contrast, the number of iterations may be dynamic if the set of operations are implemented in a “while” loop or in a “for” loop with break conditions, among other possibilities.

(18) Many processors, such as single instruction multiple data (SIMD) processors, do not have a runtime ability to control a sequence length, and as such do not directly support the execution of neural networks with dynamic lengths. However, several techniques have been introduced to adapt these processors to accommodate certain types of dynamic length neural networks. For example, for execution of simple RNNs, such as unidirectional RNNs, one solution has been to execute a maximum number of iterations, regardless of the actual length of the input, and to zero pad the input for each iteration between the actual length of the input and the maximum length. Because the output may also have the maximum length, values of the output between the actual length and the maximum length of the output are ignored by consumers of the output.

(19) While such a solution has been useful for adapting data-flow centric processors for execution of unidirectional RNNs, bidirectional RNNs have more complex structures and pose different problems. Bidirectional RNNs typically consist of two unidirectional RNNs arranged in opposite directions, one in the forward direction (e.g., positive time direction) and one in the reverse direction (e.g., negative time direction). The outputs from each of the oppositely arranged unidirectional RNNs are combined to form an overall output. Due to the arrangement of the bidirectional RNN, zero padding one or both of the inputs fed into the unidirectional RNNs will cause the bidirectional RNN to render an erroneous output. As such, new methods, systems, and other technique are needed.

(20) Embodiments described herein overcome the above-noted inability of data-flow centric processors to execute bidirectional RNNs by providing techniques for modifying (e.g., realigning) the tensors fed into and/or produced by the bidirectional RNNs at various stages/layers within the bidirectional RNNs. Such modifications may be considered a modification of the bidirectional RNN itself with additional layers, or a modification of the data being operated on by the bidirectional RNN. In some embodiments, two transformation layers are added to the bidirectional RNN structure. The first transformation layer may be added to the input that feeds the reverse unidirectional RNN. The first transformation layer may reverse the order of a first set of elements of the input and zero pad the input to a maximum length. The second transformation layer may be added to the output of the reverse unidirectional RNN before the output combiner. The second transformation layer may reverse the order of a second set of elements of the output and overwrite the output values after the second set of elements with zeros.

(21) In the following description, various examples will be described. For purposes of explanation, specific configurations and details are set forth in order to provide a thorough understanding of the examples. However, it will also be apparent to one skilled in the art that the example may be practiced without the specific details. Furthermore, well-known features may be omitted or simplified in order not to obscure the embodiments being described.

(22) FIG. 1 illustrates an example of a computational flow model **100** for a neural network. Neural networks take inspiration from the mechanics of the operation of the human brain. According to various models of the brain, the main computational element of the brain is the neuron. Neurons are

connected together with a number of elements, with elements entering a neuron being referred to as dendrites and an element leaving a neuron being referred to as an axon. A neuron accepts signals via dendrites, performs a computation on the signals, and outputs a signal on an axon. The input and output signals are referred to as activations. The axon of one neuron can branch out and be connected to the dendrites of multiple neurons. The connection between a branch of an axon and a dendrite is called a synapse.

(23) A synapse can scale the signal crossing the synapse. The scaling factor is referred to as a weight, and is thought of as the way a brain is able to learn: different weights result from different responses to input. Learning can change the weights, but the organization of the neurons and synapses need not change to obtain the learning. The static structure of the brain can thus be used as a model for a program, and the weights can reflect tasks that the program has learned to perform.

(24) Neural networks operate on the notion that a neuron's computation involves a weighted sum of input values. These weighted sums correspond to the value scaling performed by the synapses and the combining of those values in the neuron. A functional operation is performed in the neuron on the combined inputs. In the brain model, the operation appears to be a non-linear function that causes the neuron to generate an output only when the inputs cross some threshold. Thus, by analogy, the nodes of a neural network can apply a non-linear function to the weighted sum of the values input into the nodes.

(25) In the illustrated example, the model **100** includes an input layer **104**, a middle layer that is often referred to as a hidden layer **106**, and an output layer **108**. Each layer includes some number of nodes **102**. In this example, the nodes **102** of the input layer **104** are connected to each node **102** of the hidden layer **106**. The connections, which would be referred to as synapses in the brain model, are referred to as weights **110**. Also in this example, each node **102** of the hidden layer **106** has a connection or weight **110** with each node **102** of the output layer. The input layer **104** can receive inputs and can propagate the inputs to the hidden layer **106**. A neural network implementation can include multiple hidden layers. Weighted sums computed by the hidden layer **106** (or multiple hidden layers) are propagated to the output layer **108**, which can present final outputs to a user. The outputs of the nodes **102** can be referred to as activations, in keeping with the brain model.

(26) An example of a computation that can occur at each layer in the example model **100** is as follows:

$$(27) y_j = f\left(\sum_{i=1}^3 W_{ij} \times x_i + b\right)$$

(28) In the above equation, W_{ij} is a weight, x_i is an input activation, y_j is an output activation, $f()$ is a non-linear function, and b is a bias term. Various non-linear functions can be used to achieve different purposes.

(29) The model **100** can be referred to as a directed, weighted graph. In a directed graph, each connection to or from a node indicates a direction (e.g., into the node or away from the node). In a weighted graph, each connection can have a weight. Tools for developing neural networks can visualize the neural network as a directed, weighted graph, for ease of understanding and debuggability. In some cases, these tools can also be used to train the neural network and output trained weight values. Executing the neural network is then a matter of using the weights to conduct computations on input data.

(30) Neural networks with many layers can be capable of learning high-level features having more complexity and abstraction than shallower networks. As an example, a neural network can be taught to recognize images. In this example, pixels of an image can be fed into the input layer of the neural network, and the outputs of the first layer can indicate the presence of low-level features in the image, such as lines and edges. At subsequent layers, these features can be combined to measure the likely presence of higher level features: the lines can be combined into shapes, which can be further combined into sets of shapes. Given all this information, the neural network can

output a probability that the high-level features represent a particular object or scene. For example, the neural network can output whether an image contains a cat or does not contain a cat.

(31) The learning phase of a neural network is referred to as training the neural network. During training, the neural network is taught to perform a task. In learning the task, values for the weights (and possibly also the bias) are determined. The underlying program for the neural network (e.g., the organization of nodes into layers, the connections between the nodes of each layer, and the computation executed by each node), does not need to change during training. Once trained, the neural network can perform the task by computing a result using the weight values that were determined during training. For example, the neural network can output the probability that an image contains a particular object, can output the probability that an audio sequence contains a particular word, can generate a bounding box around an object in an image, or can propose an action that should be taken, etc. Running the program for the neural network is referred to as inference.

(32) There are multiple ways in which weights can be trained. One method is called supervised learning. In supervised learning, all training samples are labeled, so that inputting each training sample into a neural network produces a known result. Another method is called unsupervised learning, where the training samples are not labeled and training aims to find a structure in the data or clusters in the data. Semi-supervised learning falls between supervised and unsupervised learning. In semi-supervised learning, a subset of training data is labeled. The unlabeled data can be used to define cluster boundaries and the labeled data can be used to label the clusters.

(33) Neural networks have been used for a variety of applications, including, for example, in the areas of image and video, speech and language, medicine, game play, and robotics. In image and video, neural networks have been used for image classification, object localization and detection, image segmentation, and action recognition. In speech and language, neural networks have been used for speech recognition, machine translation, natural language processing, and audio generation. In the medical field, neural networks have been used in genomics and medical imaging. In game play, neural networks have been used to play video and board games, including games with immense numbers of possible moves such as Go. In robotics, neural networks have been used for motion planning of a robot, visual navigation, control stabilization, and driving strategies for autonomous vehicles.

(34) FIG. 2 illustrates an example of a unidirectional RNN layer **200**. The RNN layer **200** may correspond to one or more operations that may be performed on an input over multiple iterations at multiple time steps. For example, the RNN layer **200** may be unrolled into multiple RNN cells **204** that each correspond to the same one or more operations. The input to the RNN layer **200** may include N time-sequenced values from Input [0] to Input [N-1] and the output may also include N time-sequenced values from Output [0] to Output [N-1].

(35) Each of the RNN cells **204** may receive an input value along with an input state. For the first iteration, a RNN cell **204** receives the input value Input [0] and an initial state and produces the output value Output [0] and the state State [0] based on Input [0] and the initial state. Similarly, for the second iteration, a RNN cell **204** receives the input value Input [1] and the state State [0] and produces the output value Output [1] and the state State [1] based on Input [1] and State [0]. Similarly, for the N^{sup}.th iteration, a RNN cell **204** receives the input value Input [N-1] and the state State [N-2] and produces the output value Output [N-1] based on Input [N-1] and State [N-2].

(36) FIG. 3 illustrates an example of a bidirectional RNN layer **300**. The bidirectional RNN layer **300** may correspond to one or more operations that may be performed on an input over multiple iterations at multiple time steps. The bidirectional RNN layer **300** may be unrolled into multiple RNN cells **304** that each correspond to the same one or more operations. The bidirectional RNN layer **300** is an enhancement to the unidirectional RNN layer **200** and essentially consists of two unidirectional RNN layers, one operating from left to right, and the other operating from right to left. Each of the layers receives an overall input **310**, and the outputs of the two layers are merged

to produce an overall output **320**. The overall input **310** may include N time-sequenced values from Input [0] to Input [N-1] and the overall output **320** may also include N time-sequenced values from Output [0] to Output [N-1].

(37) In some embodiments, an input **312** provided to the left-to-right layer is equal to the overall input **310**, and an input **314** provided to the right-to-left layer is also equal to the overall input **310**. For the first iteration, a bidirectional RNN cell **304** of the left-to-right layer receives the input value Input [0] and an initial state and produces the output value LR Output [0] and the state LR State [0] based on Input [0] and the initial state. Also for the first iteration, a bidirectional RNN cell **304** of the right-to-left layer receives the input value Input [0] and the state RL State [1] and produces the output value RL Output [N-1] and the state RL State [0] based on Input [0] and RL State [1].

(38) For the second iteration, a bidirectional RNN cell **304** of the left-to-right layer receives the input value Input [1] and the state LR State [0] and produces the output value LR Output [1] and the state LR State [1] based on Input [1] and LR State [0]. Also for the second iteration, a bidirectional RNN cell **304** of the right-to-left layer receives the input value Input [1] and the state RL State [2] and produces the output value RL Output [N-2] and the state RL State [1] based on Input [1] and RL State [2].

(39) For the N.sup.th iteration, a bidirectional RNN cell **304** of the left-to-right layer receives the input value Input [N-1] and the state LR State [N-2] and produces the output value LR Output [N-1] and the state LR State [N-1] based on Input [N-1] and LR State [N-2]. Also for the N.sup.th iteration, a bidirectional RNN cell **304** of the right-to-left layer receives the input value Input [N-1] and an initial state and produces the output value RL Output [0] and the state RL State [N-1] based on Input [N-1] and the initial state.

(40) In some embodiments, an output **322** of the left-to-right layer includes the N time-sequenced values from LR Output [0] to LR Output [N-1] and an output **324** of the right-to-left layer includes the N time-sequenced values from RL Output [N-1] to LR Output [0]. The outputs **322** and **324** are fed into a combiner **330** which combines them by, for example, concatenating LR Output [0] with RL Output [N-1], LR Output [1] with RL Output [N-2], LR Output [N-1] with RL Output [0], and the like. The combination of outputs **322** and **324** forms the overall output **320**.

(41) FIGS. **4A** and **4B** illustrate tables that show the input states, input values, output states and output values for the bidirectional RNN layer **300**. Specifically, FIG. **4A** illustrates a table **400A** corresponding to the left-to-right layer and FIG. **4B** illustrates a table **400B** corresponding to the right-to-left layer. With respect to the notation used in FIGS. **4A** and **4B**, "Iteration=0" corresponds to the first iteration, "Iteration=1" corresponds to the second iteration, and "Iteration=N-1" corresponds to the N.sup.th iteration.

(42) When the bidirectional RNN layer **300** is statically compiled for a processor that does not support dynamic control flow, it may be compiled for a fixed length since the compiler does not have knowledge of N, which is determined at runtime. The fixed length is referred to as the maximum length, which is also referred to herein as the value "L". When the maximum length L is equal to N, the bidirectional RNN layer **300** functions properly. However, when the maximum length L is greater than N, padding the input **310** (and similarly the inputs **312** and **314**) causes an erroneous output, as demonstrated in reference to FIGS. **5A** and **5B**.

(43) FIGS. **5A** and **5B** illustrate tables that show the input states, input values, output states and output values for the bidirectional RNN layer **300**. FIG. **5A** illustrates a table **500A** showing the correct and expected behavior when the maximum length L is greater than N. In the illustrated embodiment, L=3 and N=2, and when the inference is performed for N=2 at runtime, the input for Iteration=2 is zero padded. The shaded portions of the table **500A** indicate that the values/states are irrelevant.

(44) FIG. **5B** illustrates a table **500B** showing the actual and erroneous behavior that occurs for the same scenario described in FIG. **5A** (L=3 and N=2). When the bidirectional RNN layer **300** is compiled for a static value of 3 iterations, it will execute all 3 time steps as shown in FIG. **5B**. As

indicated by arrow **502**, RL State [2] is passed to evaluate Iteration=1 instead of the initial RL state. This causes an erroneous value for RL Output [1].

(45) FIG. **6** illustrates an example of a bidirectional RNN layer **600** having additional transformation layers for modifying the input and output data of the right-to-left layer. The bidirectional RNN layer **600** may correspond to one or more operations that may be performed on an input over multiple iterations at multiple time steps. The bidirectional RNN layer **600** may be unrolled into multiple RNN cells **604** that each correspond to the same one or more operations. The bidirectional RNN layer **600** consists of two unidirectional RNN layers, one operating from left to right (alternatively referred to as the forward network or the forward RNN), and the other operating from right to left (alternatively referred to as the reverse network or the reverse RNN).

(46) An overall input **610** is received by the bidirectional RNN layer **600** and is used for an input **612** of the left-to-right layer. For the right-to-left layer, the bidirectional RNN layer **600** may include a first transformation layer **631** that performs a first transformation on the overall input **610** to produce a transformed input **614**, which is fed into the right-to-left layer. The first transformation layer **631** may reverse the order of the first N values of the overall input **610** and optionally set (e.g., overwrites or appends) the last L-N values to zero. The input **612** may include L time-sequenced values from Input [0] to Input [L-1] and the transformed input **614** may include L time-sequenced values from RL Input [0] to RL Input [L-1].

(47) The left-to-right layer produces an output **622** based on the input **612** and the right-to-left layer produces an output **624** based on the transformed input **614**. The output **622** may include L time-sequenced values from LR Output [0] to LR Output [L-1] and the output **624** may include L time-sequenced values from RL Output [0] to RL Output [L-1]. The bidirectional RNN layer **600** may include a second transformation layer **632** that performs a second transformation on the output **624** to produce a transformed output **626**, which is fed into a combiner **630**. The second transformation layer **632** may reverse the order of the first N values of the output **624** and optionally delete or overwrite the last L-N values with zeros.

(48) For the first iteration, a bidirectional RNN cell **604** of the left-to-right layer receives the input value Input [0] and an initial state and produces the output value LR Output [0] and the state LR State [0] based on Input [0] and the initial state. Also for the first iteration, a bidirectional RNN cell **604** of the right-to-left layer receives the input value RL Input [0] and an initial state and produces the output value RL Output [0] and the state RL State [0] based on RL Input [0] and the initial state.

(49) For the second iteration, a bidirectional RNN cell **604** of the left-to-right layer receives the input value Input [1] and the state LR State [0] and produces the output value LR Output [1] and the state LR State [1] based on Input [1] and LR State [0]. Also for the second iteration, a bidirectional RNN cell **604** of the right-to-left layer receives the input value RL Input [1] and the state RL State [0] and produces the output value RL Output [1] and the state RL State [1] based on RL Input [1] and RL State [0].

(50) For the Lth iteration, a bidirectional RNN cell **304** of the left-to-right layer receives the input value Input [L-1] and the state LR State [L-2] and produces the output value LR Output [L-1] and the state LR State [L-1] based on Input [L-1] and LR State [L-2]. Also for the Lth iteration, a bidirectional RNN cell **604** of the right-to-left layer receives the input value RL Input [L-1] and the state RL State [L-2] and produces the output value RL Output [L-1] and the state RL State [L-1] based on RL Input [L-1] and RL State [L-2].

(51) In some embodiments, the right-to-left layer is transformed, resulting in a transformed right-to-left layer. As shown in the illustrated embodiment, the transformed right-to-left layer may process input values from left to right and may pass states from left to right, similar to that performed by the left to right layer.

(52) The outputs **622** and **626** are fed into the combiner **630** which combines them by, for example, concatenating aligned values of the outputs, similar to that described in reference to FIG. **3**. The combination of outputs **622** and **624** forms the overall output **620**.

(53) FIGS. 7A and 7B illustrate tables that show the input states, input values, output states and output values for the bidirectional RNN layer **600**. FIG. 7A illustrates a table **700A** showing the correct and expected behavior when $N=L$.

(54) FIG. 7B illustrates a table **700B** showing the erroneous behavior (“Merged Output”) and the corrected behavior (“Corrected Merged Output”) when $N<L$ and the first transformation layer **631** and the second transformation layer **632** are either not used (for the “Merged Output”) or used (for the “Corrected Merged Output”). The shaded portions of the table **700B** indicate that the values are irrelevant. As can be observed, by using the transformation layers, RL Output [0] is moved from the end of the output values such that it becomes aligned with LR Output [1], and RL Output [1] is moved to the beginning to become aligned with LR Output [0].

(55) FIG. 8 illustrates a table **800** that shows the input states, input values, output states and output values for the right-to-left layer of the bidirectional RNN layer **600** when $L=3$ and $N=2$. In contrast to FIG. 5B, the initial RL state and RL Input [0] are processed during the same iteration, resulting in a correct value for RL Output [0]. Additionally, RL State [0] and RL Input [1] are processed during the same iteration, resulting in a correct value for RL Output [1] due to RL State [0] being the correct state and further due to RL State [0] and RL Input [1] being processed during the same iteration. The shaded portions of the table **800** indicate that the values/states are irrelevant.

(56) FIG. 9 illustrates an example of a possible first transformation performed by the first transformation layer **631** on the overall input **610** to produce the transformed input **614**. In the illustrated example, the order of the first N values of the overall input **610** is reversed such that Input [0] is switched with Input [$N-1$], Input [1] is switched with Input [$N-2$], and the like. The last $L-N$ values of the overall input **610** may also be overwritten with zeros during the first transformation.

(57) FIG. 10 illustrates an example of a possible second transformation performed by the second transformation layer **632** on the output **624** to produce the transformed output **616**. In the illustrated example, the order of the first N values of the output **624** is reversed such that RL Output [0] is switched with RL Output [$N-1$], RL Output [1] is switched with RL Output [$N-2$], and the like. The last $L-N$ values of the output **624** may also be overwritten with zeros during the second transformation. In some embodiments, the last $L-N$ values of the output **624** may also be overwritten with zeros during the second transformation.

(58) FIG. 11 includes a flowchart illustrating an example of a method **1100** for executing a bidirectional network on a data-flow centric processor. In some embodiments, the method **1100** may be considered a method for aligning tensors for execution of a bidirectional RNN on a data-flow centric processor. One or more steps of the method **1100** may be performed in a different order than the illustrated example, and one or more steps of the method **1100** may be omitted during the performance of the method **1100**. One or more steps of the method **1100** may be performed by a compiler running on a host processor, or by a data-flow centric processor such as a data-flow centric accelerator, or a runtime program running on the host processor.

(59) At step **1102**, the bidirectional network is compiled for a user-specified length (e.g., L). In some embodiments, the bidirectional network is a bidirectional RNN. In some embodiments, the bidirectional network includes a bidirectional network layer or a bidirectional RNN layer. In some embodiments, the bidirectional network includes a forward network and a reverse network, where the forward network is operable to perform a set of operations on a time-sequenced input in a forward direction and the reverse network is operable to perform a set of operations on a time-sequenced input in a reverse direction.

(60) At step **1104**, a number of runtime steps (e.g., N) is determined. In some embodiments, the number of runtime steps is determined based on an input to the bidirectional network. For example, the number of runtime steps may be determined based on the number of time steps of the input or the length of the input. For example, if the number of time-sequenced values of the input is equal to N , then the number of runtime steps may also be set equal to N . In some embodiments, the number

of runtime steps is determined while performing an inference.

(61) At step **1106**, an input is provided to the bidirectional network. In some embodiments, the input includes N time-sequenced values. In some embodiments, the input is zero padded to increase its length from N to L. In some embodiments, the input is an input tensor.

(62) At step **1108**, a forward output is generated using the forward network based on the input. The forward output may be generated when the input is provided to the forward network.

(63) At step **1110**, a first transformation is performed on the input to generate a transformed input. The first transformation may be performed by a first transformation layer of the bidirectional network. In some embodiments, performing the first transformation includes modifying an order of a set of values of the input. In some embodiments, the set of values of the input may be the first N values of the input. Modifying the order of the set of values of the input may include reversing (e.g., transposing) the order of the set of values of the input. In some embodiments, performing the first transformation includes overwriting a last set of values of the input with zeros. In some embodiments, the last set of values of the input may be the last L-N set of values. In some embodiments, the last set of values of the input may be all the values of the input after the set of values of the input.

(64) In some embodiments, reversing the order of the set of values may be implemented by the compiler, allowing the transformation to be addressed as a memory lookup instead of copying across memories and variables. In some embodiments, the compiler may identify the bidirectional network and, in response, may add the first transformation layer to the bidirectional network.

(65) At step **1112**, a reverse output is generated using the reverse network based on the transformed input. The reverse output may be generated when the transformed input is provided to the reverse network. In some embodiments, the reverse network is a transformed reverse network.

(66) At step **1114**, a second transformation is performed on the reverse output to generate a transformed output. The second transformation may be performed by a second transformation layer of the bidirectional network. In some embodiments, performing the second transformation includes modifying an order of a set of values of the reverse output. In some embodiments, the set of values of the reverse output may be the first N values of the reverse output. Modifying the order of the set of values of the reverse output may include reversing (e.g., transposing) the order of the set of values of the reverse output. In some embodiments, performing the second transformation includes overwriting a last set of values of the reverse output with zeros. In some embodiments, performing the second transformation includes deleting the last set of values of the reverse output. In some embodiments, the last set of values of the reverse output may be the last L-N set of values. In some embodiments, the last set of values of the reverse output may be all the values of the reverse output after the set of values of the reverse output.

(67) In some embodiments, reversing the order of the set of values may be implemented by the compiler, allowing the transformation to be addressed as a memory lookup instead of copying across memories and variables. In some embodiments, the compiler may identify the bidirectional network and, in response, may add the first transformation layer to the bidirectional network.

(68) At step **1116**, the forward output and the transformed output are combined to produce an output or an overall output. The forward output and the transformed output may be combined by, for example, concatenating corresponding values such as a first value of the forward output with a first value of the transformed output, a second value of the forward output with a second value of the transformed output, and the like.

(69) FIG. 12 includes a block diagram illustrating an example of a host system **1200** on which a compiler, a data-flow centric processor, and a host processor, such as are described herein, can run. The illustrated host system **1200** is an example of a computing device, and includes a processor **1202**, a processor memory **1204**, at least one storage device **1206**, various Input/Output (I/O) devices **1208**, and at least one network interface **1210**. In the example of FIG. 12, the host system **1200** also includes an acceleration engine **1212**, which is an integrated circuit device that can

accelerate certain operations or computations performed by the host system **1200**. In various examples, the host system **1200** can be implemented as a server in a data center, a desktop computer, a laptop computer, a tablet computer, or a smartphone, among other examples. In some examples, operations or components discussed below as performed or included in the host system **1200** can be performed or included in other computer devices. For example, the compiler **1230** can execute on the host system **1200** while the acceleration engine **1212** is located at a different host system.

(70) The processor **1202** is an integrated circuit device that can execute program code, in the form of instructions. The program code can be used for various software applications or tools, such as an operating system **1220**, the debugger **1246**, or the illustrated compiler **1230**. While the processor **1202** is executing a program, the instructions for the program can be stored in the processor memory **1204**. The instructions can also be stored elsewhere, such as on the storage device **1206**, and can be loaded into the processor memory **1204** when needed by the processor **1202**. The processor **1202** can also use the processor memory **1204** for temporary storage of other data on which the processor **1202** is operating. In various examples, the processor memory **1204** is a volatile memory type, such as a type of Random Access Memory, though non-volatile memory types can, alternatively or additionally, be used for the processor memory **1204**.

(71) The storage device **1206** is an example of a device that can include non-volatile memory. For example, the storage device **1206** can be a magnetic disk drive, a solid state drive, or an optical drive, among other examples. The storage device **1206** can further be non-transitory, such that program code and other data stored on the storage device **1206** remains present when the storage device **1206** is not powered on.

(72) The storage device **1206** is one example of a peripheral device, which are components that can be coupled to the host system **1200** to add functionality to the host system **1200**. Other examples of peripheral devices include the Input/Output devices **1208** and the network interface **1210**. The Input/Output devices **1208** can include user input and output devices, such as keyboards, mice, touch screens, microphones, display screens, speakers, printers, and scanners, among other examples. The network interface **1210**, which can be implemented using a network interface card, can provide access to one or more networks. The network interface **1210** can include, for example, a physical port for connecting a network cable and/or wireless antennas for communicating with Wi-Fi and/or cellular networks. The network interface **1210** can also be described as an I/O device.

(73) The acceleration engine **1212** is also another type of peripheral device or I/O device. The acceleration engine **1212** is a device that is purpose built to perform certain operations that can be performed by the processor **1202**, but can be performed faster by the acceleration engine **1212**. For example, the acceleration engine **1212** can be a neural network accelerator, and, as such, may be able to perform the large scale, parallel computations of a neural network more efficiently than when the computations are performed by the processor **1202**. As another example, the acceleration engine **1212** can be a GPU, and may be optimized to perform the computations needed for graphics rendering. Other examples of devices that can be implemented by the acceleration engine **1212** include cryptographic accelerators, compression and decompression accelerators, 3-D accelerators, regular expression accelerators, security accelerators, and others.

(74) In various examples, the acceleration engine **1212** can execute program code to perform certain operations. For example, when the acceleration engine **1212** is a neural network accelerator, the acceleration engine **1212** can be programmed to execute a particular neural network, such as one that performs image recognition or one that performs machine translation. As a further example, to support the execution of a neural network, the acceleration engine **1212** can be programmed to perform operations such as copying data for the neural network from processor memory **1204** (for example) into the acceleration engine **1212**, copying input data for the neural network from processor memory **1204** into the acceleration engine **1212**, and/or copying results from the acceleration engine **1212** into the processor memory **1204**, among other examples.

(75) To generate program code for the acceleration engine **1212**, in various examples, the host system **1200** can execute the compiler **1230**. Compilers, in general, are software programs that translate program code written in a human-readable language into a format (e.g., machine instructions) that can be read and processed by an integrated circuit device. In the example of FIG. **12**, the acceleration engine **1212** is a neural network accelerator and the compiler **1230** is for compiling a neural network description into instructions to be executed on the acceleration engine **1212**. When the acceleration engine **1212** implements a different type of accelerator, another compiler can be used.

(76) The compiler **1230** can be activated, for example, when the operating system **1220** receives keyboard, mouse, touchscreen, voice commands, or other inputs from the Input/Output devices **1208**. The inputs can further include parameters for the compiler **1230**, such as the input code **1242** to compile and configuration options for the compilation process. Once the compiler **1230** is activated, the processor **1202** can load the instructions for the compiler **1230** into the processor memory **1204**, and can execute the instructions.

(77) In the example of FIG. **12**, the compiler **1230** includes a first stage **1232**, a second stage **1236**, and a third stage **1240**, which each perform different operations to produce compiled code **1244**. In other examples, the compiler **1230** can combine the operations of the first stage **1232**, second stage **1236**, and/or third stage **1240** into fewer stages, or can divide the operations of one of the stages into multiple stages.

(78) The first stage **1232** can receive and process input code **1242**. The input code **1242** can describe a program in a high-level programming language, such as Java, C++, or Tensorflow, among many other examples. The input code **1242** can describe, for example, steps to perform image recognition, speech recognition, machine translation, or other operations. The input code **1242** can be obtained, for example, from the storage device **1206**. Alternatively, though not illustrated here, the input code **1242** may be located in the processor memory **1204** or can be obtained from a network location, using the network interface **1210**. Processing of the input code **1242** can include sorting the operations described in the input code **1242** into layers, where the outputs of one layer provide the inputs to a next layer. Processing can also include identifying steps to be performed by the processor **1202**, rather than by the acceleration engine **1212**. For example, the processor **1202**, through the execution of a driver **1222**, may need to perform steps such as configuring Direct Memory Access (DMA) descriptors for moving data into or out of the acceleration engine **1212**, among other examples.

(79) The output **1234** of the first stage **1232** can be organized, for example, in the layers, nodes, and connections between nodes of a neural network. The second stage **1236** can perform intermediate processing on this output **1234**. For example, the operations performed in any one layer, or at any one node in a layer, may be too many for the acceleration engine **1212** to perform at the same time. The acceleration engine **1212** may, for example, have a limited amount of local storage space for the data needed for a computation, or the computations may be more than the acceleration engine **1212** can perform at one time. In this example, the first stage **1232** can break the operations of the layer or node down into smaller operations, which can fit into the acceleration engine's local memory and/or can fit into the computing capacity of the acceleration engine **1212**. Processing of the output **1234** of the first stage **1232** can include other steps, such as scheduling, or determining the order in which the acceleration engine **1212** and/or processor **1202** will perform operations, among other examples.

(80) In various examples, the output **1238** of the second stage **1236** includes the various steps to be performed by components of the acceleration engine **1212**, in the order that the steps are to be performed. The output **1238** can be represented, for example, as a data flow graph, where the nodes in the graph represent memory operations, computations, and other operations, and the edges or connections between the nodes represent dependencies between the nodes, such as data dependencies, memory dependencies, or operational dependencies, among other examples.

(81) The third stage **1240** can operate on the output **1238** of the second stage **1236**, and perform various steps before producing the instructions that are to be executed by the acceleration engine **1212**. These steps can include, for example, removing redundant dependencies, resolving or handling dependencies between nodes by inserting synchronization instructions into the code, identifying possible optimizations in memory usage or memory bandwidth usage, and other operations.

(82) The output of the third stage **1240** is compiled code **1244**, which may include machine instructions in binary format. In some examples, the compiled code **1244** can be stored in the processor memory **1204**. Alternatively or additionally, the compiled code **1244** can be copied to the storage device **1206** or to a network location. As noted above, the acceleration engine **1212** may be located at a different host system, in which case the compiled code **1244** can be sent over the network interface **1210** to the other host system.

(83) In the example of FIG. 12, the host system **1200** can execute a driver **1222**, which can also be referred to as a device driver or runtime driver, that manages the acceleration engine **1212**. The driver **1222** can provide an interface between applications executing on the host system **1200** (or on another host system) and the acceleration engine **1212**. For example, the driver **1222** can provide an Application Program Interface (API) that defines functions for feeding input data to the acceleration engine **1212** and defining the operation to perform on the input data. In this and other examples, the driver **1222** can configure the acceleration engine **1212** to perform the operation. For example, the driver **1222** can identify a neural network that the acceleration engine **1212** is to execute, as well as the location in the processor memory **1204** or on the storage device **1206** where the compiled code **1244** for the neural network is located. The driver **1222** can further load into the acceleration engine **1212** or cause the acceleration engine **1212** to load the compiled code **1244**, can load or cause the acceleration engine **1212** to load the input data on which the neural network is to operate, and/or can cause the acceleration engine **1212** to begin executing on the input data. Once the acceleration engine **1212** has finished, the acceleration engine **1212** can notify the driver **1222**, and the driver **1222** can deliver a result back to the application that requested the result.

(84) FIG. 13 includes a block diagram that illustrates an example of an acceleration engine **1300**. The acceleration engine **1300** is an example of an integrated circuit that can include one or more accelerators **1302a-1302n** that may be similar to the accelerator illustrated in FIG. 14.

(85) In the example of FIG. 13, the acceleration engine **1300** includes multiple accelerators **1302a-1302n**, each of which can perform a set of operations. In various examples, the accelerators **1302a-1302n** are for particular types of operations, so that the accelerators **1302a-1302n** can perform the operations much faster than when similar operations are performed by a general purpose processor. In various examples, to perform a set of operations, input data on which the operations are to be performed must first be moved into the accelerators **1302a-1302n**. Additionally, in some cases, program code is also moved into the accelerators **1302a-1302n**, which programs the operations that the accelerators **1302a-1302n** will perform on the data. In the illustrated example, the acceleration engine **1300** includes *n* accelerators **1302a-1302n**. Examples of accelerators that can be included in the acceleration engine **1300** include graphics accelerators, floating point accelerators, neural network accelerators, and others. In various examples, the accelerators **1302a-1302n** can each be the same (e.g., each of the is a graphics accelerator) or can be different (e.g., the accelerators **1302a-1302n** include a graphics accelerator, a floating point accelerator, and neural network accelerator).

(86) The example acceleration engine **1300** further includes DRAM controllers **1342a-1342k** for communicating with an external memory. The external memory is implemented, in this example, using DRAM **1330**. In the illustrated example, the acceleration engine **1300** includes *k* DRAM controllers **1342a-1342k**, each of which may be able to communicate with an independent set of banks of DRAM. In other examples, other types of RAM technology can be used for the external memory. The DRAM controllers **1342a-1342k** can also be referred to as memory controllers.

(87) In various examples, input data and/or program code for the accelerators **1302a-1302n** can be stored in the DRAM **1330**. Different programs can cause the accelerators **1302a-1302n** to perform different operations. For example, when one of the accelerators is a neural network accelerator, one program can configure the neural network accelerator to perform speech recognition while another program can configure the neural network accelerator to perform image recognition. In various examples, different accelerators **1302a-1302n** can be programmed with different programs, so that each performs a different set of operations. In various examples, the processors **1348a-1348s** can manage moving of program code from the DRAM **1330** to the accelerators **1302a-1302n**.

(88) The example acceleration engine **1300** further includes I/O controllers **1344a-1344p** for communicating with I/O devices **1332** in the system. The acceleration engine **1300** can communicate with I/O devices over, for example, a processor bus. In some examples, the processor bus can be implemented using Peripheral Component Interconnect (PCI) and/or a variation of the PCI bus protocol. The processor bus can connect the acceleration engine **1300** to I/O devices such as, for example, input and output devices, memory controllers, storage devices, and/or network interface cards, among other things. In some examples, the I/O controllers **1344a-1344p** can enable the acceleration engine **1300** to act as an I/O device for a host processor. For example, the acceleration engine **1300** can be the recipient of input data from the host processor, and a command indicating an operation to be performed on the input data (e.g., a particular computation or analysis). In the illustrated example, the acceleration engine **1300** includes p I/O controllers **1344a-1344p**, each of which may include a separate root complex and may communicate with a separate set of I/O devices **1332**. In other examples, other standardized bus protocols, such as Ultra Path Interconnect (UPI) can be used for the host bus. In other examples, a proprietary bus protocol can be used.

(89) Movement of data in the acceleration engine **1300** can be managed by one or more processors **1348a-1348s**, which can also be referred to as data management processors. In the example of FIG. 13, the acceleration engine **1300** includes s processors **1348a-1348s** incorporated into the device (e.g., on the same silicon die). In other examples, the processors **1348a-1348s** can be external to the acceleration engine **1300** (e.g., on a different die and/or in a different package). In some examples, the processors **1348a-1348s** can manage the movement of data from I/O devices **1332** to the accelerators **1302a-1302n** or the DRAM **1330**. For example, input data may be located at an I/O device **1332** or in processor memory, and the processors **1348a-1348s** can move the input from the I/O device **1332** or processor memory into an accelerator or into DRAM **1330**. As another example, program code for the accelerators **1302a-1302n** may be located on an I/O device **1332** or in processor memory.

(90) The example acceleration engine **1300** further includes DMA engines **1346a-1346d** that can move data between the accelerators **1302a-1302n**, DRAM controllers **1342a-1342k**, and I/O controllers **1344a-1344p**. In the illustrated example, the acceleration engine **1300** includes d DMA engines **1346a-1346d**. In some implementations, the DMA engines **1346a-1346d** can be assigned to specific tasks, such as moving data from the DRAM controllers **1342a-1342d** to the accelerators **1302a-1302n**, or moving data between the I/O controllers **1344a-1344p** and the accelerators **1302a-1302n**. These tasks can be assigned, for example, by enqueueing descriptors with the DMA engines **1346a-1346d**, where a descriptor identifies an address for a block of data and an operation (e.g., a read or a write) to perform. A descriptor, for example, can direct a DMA engine to instruct a DMA controller to read a block of data from DRAM **1330**. A descriptor can, as a further example, instruct the DMA engine to write data, read by the DMA controller, to an accelerator. Further descriptors can be used to move data from an accelerator to DRAM **1330**.

(91) In various examples, each of the processors **1348a-1348s** can be responsible for managing the data movement for a different accelerator. In some examples, a processor may manage the data movement for more than one accelerator. Similarly, in various examples, each of the processors **1348a-1348s** can be assigned to one or more DMA engines **1346a-1346d**. In these and other

examples, associations between processors **1348a-1348s**, accelerators **1302a-1302n**, and DMA engines **1346a-1346d** are determined by program code being executed by each respective processor.

(92) In the example acceleration engine **1300**, the various components can communicate over a chip interconnect **1320**. The chip interconnect **1320** primarily includes wiring for routing data between the components of the acceleration engine **1300**. In some cases, the chip interconnect **1320** can include a minimal amount of logic, such as multiplexors to control the direction of data, flip-flops for handling clock domain crossings, and timing logic.

(93) FIG. **14** is a block diagram illustrating an example of an integrated circuit device that can execute a unidirectional or bidirectional RNN layer or network. The example of FIG. **14** illustrates an accelerator **1402**. In various examples, the accelerator **1402**, for a set of input data (e.g., input data **1450**), can execute computations using a processing engine array **1410**, an activation engine **1416**, and/or a pooling engine **1418**. In some examples, the example accelerator **1402** may be an integrated circuit component of a processor, such as a neural network processor. The processor may have other integrated circuit components, including additional accelerator engines.

(94) In various implementations, the memory subsystem **1404** can include multiple memory banks **1414**. In these implementations, each memory bank **1414** can be independently accessible, meaning that the read of one memory bank is not dependent on the read of another memory bank. Similarly, writing to one memory bank does not affect or limit writing to a different memory bank. In some cases, each memory bank can be read and written at the same time. Various techniques can be used to have independently accessible memory banks **1414**. For example, each memory bank can be a physically separate memory component that has an address space that is separate and independent of the address spaces of each other memory bank. In this example, each memory bank may have at least one read channel and may have at least one separate write channel that can be used at the same time. In these examples, the memory subsystem **1404** can permit simultaneous access to the read or write channels of multiple memory banks. As another example, the memory subsystem **1404** can include arbitration logic such that arbitration between, for example, the outputs of multiple memory banks **1414** can result in more than one memory bank's output being used. In these and other examples, though globally managed by the memory subsystem **1404**, each memory bank can be operated independently of any other.

(95) Having the memory banks **1414** be independently accessible can increase the efficiency of the accelerator **1402**. For example, values can be simultaneously read and provided to each row of the processing engine array **1410**, so that the entire processing engine array **1410** can be in use in one clock cycle. As another example, the memory banks **1414** can be read at the same time that results computed by the processing engine array **1410** are written to the memory subsystem **1404**. In contrast, a single memory may be able to service only one read or write at a time. With a single memory, multiple clock cycles can be required, for example, to read input data for each row of the processing engine array **1410** before the processing engine array **1410** can be started.

(96) In various implementations, the memory subsystem **1404** can be configured to simultaneously service multiple clients, including the processing engine array **1410**, the activation engine **1416**, the pooling engine **1418**, and any external clients that access the memory subsystem **1404** over a communication fabric **1420**. In some implementations, being able to service multiple clients can mean that the memory subsystem **1404** has at least as many memory banks as there are clients. In some cases, each row of the processing engine array **1410** can count as a separate client. In some cases, each column of the processing engine array **1410** can output a result, such that each column can count as a separate write client. In some cases, output from the processing engine array **1410** can be written into the memory banks **1414** that can then subsequently provide input data for the processing engine array **1410**. As another example, the activation engine **1416** and the pooling engine **1418** can include multiple execution channels, each of which can be separate memory clients. The memory banks **1414** can be implemented, for example, using static random access

memory (SRAM).

(97) In various implementations, the memory subsystem **1404** can include control logic. The control logic can, for example, keep track of the address spaces of each of the memory banks **1414**, identify memory banks **1414** to read from or write to, and/or move data between the memory banks **1414**. In some implementations, memory banks **1414** can be hardwired to particular clients. For example, a set of memory banks **1414** can be hardwired to provide values to the rows of the processing engine array **1410**, with one memory bank servicing each row. As another example, a set of memory banks can be hard wired to receive values from columns of the processing engine array **1410**, with one memory bank receiving data for each column.

(98) The processing engine array **1410** is the computation matrix of the example accelerator **1402**. The processing engine array **1410** can, for example, execute parallel integration, convolution, correlation, and/or matrix multiplication, among other things. The processing engine array **1410** includes multiple processing engines **1411**, arranged in rows and columns, such that results output by one processing engine **1411** can be input directly into another processing engine **1411**.

Processing engines **1411** that are not on the outside edges of the processing engine array **1410** thus can receive data to operate on from other processing engines **1411**, rather than from the memory subsystem **1404**.

(99) In various examples, the processing engine array **1410** uses systolic execution, in which data arrives at each processing engine **1411** from different directions at regular intervals. In some examples, input data can flow into the processing engine array **1410** from the left and weight values can be loaded at the top. In some examples weights and input data can flow from the left and partial sums can flow from top to bottom. In these and other examples, a multiply-and-accumulate operation moves through the processing engine array **1410** as a diagonal wave front, with data moving to the right and down across the array. Control signals can be input at the left at the same time as weights, and can flow across and down along with the computation.

(100) In various implementations, the number of columns in the processing engine array **1410** determines the computational capacity of the processing engine array **1410**, and the number of rows determines the required memory bandwidth for achieving maximum utilization of the processing engine array **1410**. The processing engine array **1410** can have, for example, 64 columns and 428 rows, or some other number of columns and rows.

(101) An example of a processing engine **1411** is illustrated in FIG. **14** in an inset diagram. As illustrated by this example, a processing engine **1411** can include a multiplier-accumulator circuit. Inputs from the left can include, for example, input data i and a weight value w , where the input data is a value taken from either a set of input data or a set of intermediate results, and the weight value is from a set of weight values that connect one layer of the neural network to the next. A set of input data can be, for example, an image being submitted for identification or object recognition, an audio clip being provided for speech recognition, a string of text for natural language processing or machine translation, or the current state of a game requiring analysis to determine a next move, among other things. In some examples, the input data and the weight value are output to the right, for input to the next processing engine **1411**.

(102) In the illustrated example, an input from above can include a partial sum, p_{in} , provided either from another processing engine **1411** or from a previous round of computation by the processing engine array **1410**. When starting a computation for a new set of input data, the top row of the processing engine array **1410** can receive a fixed value for p_{in} , such as zero. As illustrated by this example, i and w are multiplied together and the result is summed with p_{in} to produce a new partial sum, p_{out} , which can be input into another processing engine **1411**. Various other implementations of the processing engine **1411** are possible.

(103) Outputs from the last row in the processing engine array **1410** can be temporarily stored in the results buffer **1412**. The results can be intermediate results, which can be written to the memory banks **1414** to be provided to the processing engine array **1410** for additional computation.

Alternatively, the results can be final results, which, once written to the memory banks **1414** can be read from the memory subsystem **1404** over the communication fabric **1420**, to be output by the system.

(104) In some implementations, the accelerator **1402** includes an activation engine **1416**. In these implementations, the activation engine **1416** can combine the results from the processing engine array **1410** into one or more output activations. For example, for a convolutional neural network, convolutions from multiple channels can be summed to produce an output activation for a single channel. In other examples, accumulating results from one or more columns in the processing engine array **1410** may be needed to produce an output activation for a single node in the neural network. In some examples, activation engine **1416** can be bypassed.

(105) In various examples, the activation engine **1416** can include multiple separate execution channels. In these examples, the execution channels can correspond to the columns of the processing engine array **1410**, and can perform an operation on the outputs of a column, the result of which can be stored in the memory subsystem **1404**. In these examples, the activation engine **1416** may be able to perform between 1 and n parallel computations, where n is equal to the number of columns in the processing engine array **1410**. In some cases, one or more of the computations can be performed simultaneously. Examples of computations that each execution channel can perform include exponentials, squares, square roots, identities, binary steps, bipolar steps, sigmoidals, and ramps, among other examples.

(106) In some implementations, the accelerator **1402** can include a pooling engine **1418**. Pooling is the combining of outputs of the columns of the processing engine array **1410**. Combining can include for example, computing a maximum value, a minimum value, an average value, a median value, a summation, a multiplication, or another logical or mathematical combination. In various examples, the pooling engine **1418** can include multiple execution channels that can operating on values from corresponding columns of the processing engine array **1410**. In these examples, the pooling engine **1418** may be able to perform between 1 and n parallel computations, where n is equal to the number of columns in the processing engine array **1410**. In various examples, execution channels of the pooling engine **1418** can operate in parallel and/or simultaneously. In some examples, the pooling engine **1418** can be bypassed.

(107) Herein, the activation engine **1416** and the pooling engine **1418** may be referred to collectively as execution engines. The processing engine array **1410** is another example of an execution engine. Another example of an execution engine is a Direct Memory Access (DMA) engine, which may be located outside the accelerator **1402**.

(108) Input data **1450** can arrive over the communication fabric **1420**. The communication fabric **1420** can connect the accelerator **1402** to other components of a processor, such as a DMA engine that can obtain input data **1450** from an Input/Output (I/O) device, a storage drive, or a network interface. The input data **1450** can be, for example one-dimensional data, such as a character string or numerical sequence, or two-dimensional data, such as an array of pixel values for an image or frequency and amplitude values over time for an audio signal. In some examples, the input data **1450** can be three-dimensional, as may be the case with, for example, the situational information used by a self-driving car or virtual reality data. In some implementations, the memory subsystem **1404** can include a separate buffer for the input data **1450**. In some implementations, the input data **1450** can be stored in the memory banks **1414** when the accelerator **1402** receives the input data **1450**.

(109) In some examples, the accelerator **1402** can implement a neural network processing engine. In these examples, the accelerator **1402**, for a set of input data **1450**, can execute a neural network to perform a task for which the neural network was trained. Executing a neural network on a set of input data can be referred to as inference or performing inference.

(110) The weights for the neural network can be stored in the memory subsystem **1404**, along with input data **1450** on which the neural network will operate. The neural network can also include

instructions, which can program the processing engine array **1410** to perform various computations on the weights and the input data. The instructions can also be stored in the memory subsystem **1404**, in the memory banks **1414** or in a separate instruction buffer. The processing engine array **1410** can output intermediate results, which represent the outputs of individual layers of the neural network. In some cases, the activation engine **1416** and/or pooling engine **1418** may be enabled for computations called for by certain layers of the neural network. The accelerator **1402** can store the intermediate results in the memory subsystem **1404** for inputting into the processing engine array **1410** to compute results for the next layer of the neural network. The processing engine array **1410** can further output final results from a last layer of the neural network. The final results can be stored in the memory subsystem **1404** and then be copied out to host processor memory or to another location.

(111) FIG. **15** includes a diagram of an example network **1500**, which can include one or more host systems, such as the host system illustrated in FIG. **12**. For example, the example network **1500** of FIG. **15** includes multiple nodes **1502a-1502h**, one or more of which can be a host system such as is illustrated in FIG. **12**. Others of the nodes **1502a-1502h** can be other computing devices, each of which include at least a memory for storing program instructions, a processor for executing the instructions, and a network interface for connecting to the network **1500**.

(112) In various examples, the network **1500** can be used to process data. For example, input data can be received at one of the nodes **1502a-1502h** or from other networks **1508** with which the network **1500** can communicate. In this example, the input data can be directed to a node in the network **1500** that includes an acceleration engine, for the acceleration engine to operate on and produce a result. The result can then be transferred to the node or other network from which the input data was received. In various examples, input data can be accumulated from various sources, including one or more of the nodes **1502a-1502h** and/or computing devices located in the other networks **1508**, and the accumulated input data can be directed to one or more host systems in the network **1500**. Results from the host systems can then be distributed back to the sources from which the input data was gathered.

(113) In various examples, one or more of the nodes **1502a-1502h** can be responsible for operations such as accumulating input data for host systems to operate on, keeping track of which host systems are busy and which can accept more work, determining whether the host systems are operating correctly and/or most efficiently, monitoring network security, and/or other management operations.

(114) In the example of FIG. **15**, the nodes **1502a-1502h** are connected to one another using a switched architecture with point-to-point links. The switched architecture includes multiple switches **1504a-1504d**, which can be arranged in a multi-layered network such as a Clos network. A network device that filters and forwards packets between local area network (LAN) segments may be referred to as a switch. Switches generally operate at the data link layer (layer 2) and sometimes the network layer (layer 3) of the Open System Interconnect (OSI) Reference Model and may support several packet protocols. The switches **1504a-1504d** of FIG. **15** may be connected to the nodes **1502a-1502h** and provide multiple paths between any two nodes.

(115) The network **1500** may also include one or more network devices for connection with other networks **1508**, such as a router **1506**. Routers use headers and forwarding tables to determine the best path for forwarding the packets, and use protocols such as internet control message protocol (ICMP) to communicate with each other and configure the best route between any two devices. The router **1506** of FIG. **15** can be used to connect to other networks **1508** such as subnets, LANs, wide area networks (WANs), and/or the Internet.

(116) In some examples, network **1500** may include any one or a combination of many different types of networks, such as cable networks, the Internet, wireless networks, cellular networks and other private and/or public networks. The interconnected switches **1504a-1504d** and the router **1506**, if present, may be referred to as a switch fabric **1510**, a fabric, a network fabric, or simply a

network. In the context of a computer network, terms “fabric” and “network” may be used interchangeably herein.

(117) The nodes **1502a-1502h** may be any combination of host systems, processor nodes, storage subsystems, and I/O chassis that represent user devices, service provider computers or third party computers.

(118) User devices may include computing devices to access an application **1532** (e.g., a web browser or mobile device application). In some aspects, the application **1532** may be hosted, managed, and/or provided by a computing resources service or service provider. The application **1532** may allow the user(s) to interact with the service provider computer(s) to, for example, access web content (e.g., web pages, music, video, etc.). The user device(s) may be a computing device such as for example a mobile phone, a smart phone, a personal digital assistant (PDA), a laptop computer, a netbook computer, a desktop computer, a thin-client device, a tablet computer, an electronic book (e-book) reader, a gaming console, etc. In some examples, the user device(s) may be in communication with the service provider computer(s) via the other network(s) **1508**. Additionally, the user device(s) may be part of the distributed system managed by, controlled by, or otherwise part of the service provider computer(s) (e.g., a console device integrated with the service provider computers).

(119) The node(s) of FIG. **15** may also represent one or more service provider computers. One or more service provider computers may provide a native application that is configured to run on the user devices, which user(s) may interact with. The service provider computer(s) may, in some examples, provide computing resources such as, but not limited to, client entities, low latency data storage, durable data storage, data access, management, virtualization, cloud-based software solutions, electronic content performance management, and so on. The service provider computer(s) may also be operable to provide web hosting, databasing, computer application development and/or implementation platforms, combinations of the foregoing or the like to the user(s). In some examples, the service provider computer(s) may be provided as one or more virtual machines implemented in a hosted computing environment. The hosted computing environment may include one or more rapidly provisioned and released computing resources. These computing resources may include computing, networking and/or storage devices. A hosted computing environment may also be referred to as a cloud computing environment. The service provider computer(s) may include one or more servers, perhaps arranged in a cluster, as a server farm, or as individual servers not associated with one another and may host the application **1532** and/or cloud-based software services. These servers may be configured as part of an integrated, distributed computing environment. In some aspects, the service provider computer(s) may, additionally or alternatively, include computing devices such as for example a mobile phone, a smart phone, a personal digital assistant (PDA), a laptop computer, a desktop computer, a netbook computer, a server computer, a thin-client device, a tablet computer, a gaming console, etc. In some instances, the service provider computer(s), may communicate with one or more third party computers.

(120) In one example configuration, the node(s) **1502a-1502h** may include at least one memory **1518** and one or more processing units (or processor(s) **1520**). The processor(s) **1520** may be implemented in hardware, computer-executable instructions, firmware, or combinations thereof. Computer-executable instruction or firmware implementations of the processor(s) **1520** may include computer-executable or machine-executable instructions written in any suitable programming language to perform the various functions described.

(121) In some instances, the hardware processor(s) **1520** may be a single core processor or a multi-core processor. A multi-core processor may include multiple processing units within the same processor. In some examples, the multi-core processors may share certain resources, such as buses and second or third level caches. In some instances, each core in a single or multi-core processor may also include multiple executing logical processors (or executing threads). In such a core (e.g.,

those with multiple logical processors), several stages of the execution pipeline and also lower level caches may also be shared.

(122) The memory **1518** may store program instructions that are loadable and executable on the processor(s) **1520**, as well as data generated during the execution of these programs. Depending on the configuration and type of the node(s) **1502a-1502h**, the memory **1518** may be volatile (such as RAM) and/or non-volatile (such as ROM, flash memory, etc.). The memory **1518** may include an operating system **1528**, one or more data stores **1530**, one or more application programs **1532**, one or more drivers **1534**, and/or services for implementing the features disclosed herein.

(123) The operating system **1528** may support nodes **1502a-1502h** basic functions, such as scheduling tasks, executing applications, and/or controller peripheral devices. In some implementations, a service provider computer may host one or more virtual machines. In these implementations, each virtual machine may be configured to execute its own operating system. Examples of operating systems include Unix, Linux, Windows, Mac OS, IOS, Android, and the like. The operating system **1528** may also be a proprietary operating system.

(124) The data stores **1530** may include permanent or transitory data used and/or operated on by the operating system **1528**, application programs **1532**, or drivers **1534**. Examples of such data include web pages, video data, audio data, images, user data, and so on. The information in the data stores **1530** may, in some implementations, be provided over the network(s) **1508** to user devices. In some cases, the data stores **1530** may additionally or alternatively include stored application programs and/or drivers. Alternatively or additionally, the data stores **1530** may store standard and/or proprietary software libraries, and/or standard and/or proprietary application user interface (API) libraries. Information stored in the data stores **1530** may be machine-readable object code, source code, interpreted code, or intermediate code.

(125) The drivers **1534** include programs that may provide communication between components in a node. For example, some drivers **1534** may provide communication between the operating system **1528** and additional storage **1522**, network device **1524**, and/or I/O device **1526**. Alternatively or additionally, some drivers **1534** may provide communication between application programs **1532** and the operating system **1528**, and/or application programs **1532** and peripheral devices accessible to the service provider computer. In many cases, the drivers **1534** may include drivers that provide well-understood functionality (e.g., printer drivers, display drivers, hard disk drivers, Solid State Device drivers). In other cases, the drivers **1534** may provide proprietary or specialized functionality.

(126) The service provider computer(s) or servers may also include additional storage **1522**, which may include removable storage and/or non-removable storage. The additional storage **1522** may include magnetic storage, optical disks, solid state disks, flash memory, and/or tape storage. The additional storage **1522** may be housed in the same chassis as the node(s) **1502a-1502h** or may be in an external enclosure. The memory **1518** and/or additional storage **1522** and their associated computer-readable media may provide non-volatile storage of computer-readable instructions, data structures, program modules, and other data for the computing devices. In some implementations, the memory **1518** may include multiple different types of memory, such as SRAM, DRAM, or ROM.

(127) The memory **1518** and the additional storage **1522**, both removable and non-removable, are examples of computer-readable storage media. For example, computer-readable storage media may include volatile or non-volatile, removable or non-removable media implemented in a method or technology for storage of information, the information including, for example, computer-readable instructions, data structures, program modules, or other data. The memory **1518** and the additional storage **1522** are examples of computer storage media. Additional types of computer storage media that may be present in the node(s) **1502a-1502h** may include, but are not limited to, PRAM, SRAM, DRAM, RAM, ROM, EEPROM, flash memory or other memory technology, CD-ROM, DVD or other optical storage, magnetic cassettes, magnetic tape, magnetic disk storage or other

magnetic storage devices, solid state drives, or some other medium which can be used to store the desired information and which can be accessed by the node(s) **1502a-1502h**. Computer-readable media also includes combinations of any of the above media types, including multiple units of one media type.

(128) Alternatively or additionally, computer-readable communication media may include computer-readable instructions, program modules or other data transmitted within a data signal, such as a carrier wave or other transmission. However, as used herein, computer-readable storage media does not include computer-readable communication media.

(129) The node(s) **1502a-1502h** may also include I/O device(s) **1526**, such as a keyboard, a mouse, a pen, a voice input device, a touch input device, a display, speakers, a printer, and the like. The node(s) **1502a-1502h** may also include one or more communication channels **1536**. A communication channel **1536** may provide a medium over which the various components of the node(s) **1502a-1502h** can communicate. The communication channel or channels **1536** may take the form of a bus, a ring, a switching fabric, or a network.

(130) The node(s) **1502a-1502h** may also contain network device(s) **1524** that allow the node(s) **1502a-1502h** to communicate with a stored database, another computing device or server, user terminals and/or other devices on the network(s) **1500**.

(131) In some implementations, the network device **1524** is a peripheral device, such as a PCI-based device. In these implementations, the network device **1524** includes a PCI interface for communicating with a host device. The term “PCI” or “PCI-based” may be used to describe any protocol in the PCI family of bus protocols, including the original PCI standard, PCI-X, Accelerated Graphics Port (AGP), and PCI-Express (PCIe) or any other improvement or derived protocols that are based on the PCI protocols discussed herein. The PCI-based protocols are standard bus protocols for connecting devices, such as a local peripheral device to a host device. A standard bus protocol is a data transfer protocol for which a specification has been defined and adopted by various manufacturers. Manufacturers ensure that compliant devices are compatible with computing systems implementing the bus protocol, and vice versa. As used herein, PCI-based devices also include devices that communicate using Non-Volatile Memory Express (NVMe). NVMe is a device interface specification for accessing non-volatile storage media attached to a computing system using PCIe. For example, the bus interface module may implement NVMe, and the network device **1524** may be connected to a computing system using a PCIe interface.

(132) A PCI-based device may include one or more functions. A “function” describes operations that may be provided by the network device **1524**. Examples of functions include mass storage controllers, network controllers, display controllers, memory controllers, serial bus controllers, wireless controllers, and encryption and decryption controllers, among others. In some cases, a PCI-based device may include more than one function. For example, a PCI-based device may provide a mass storage controller and a network adapter. As another example, a PCI-based device may provide two storage controllers, to control two different storage resources. In some implementations, a PCI-based device may have up to eight functions.

(133) In some implementations, the network device **1524** may include single-root I/O virtualization (SR-IOV). SR-IOV is an extended capability that may be included in a PCI-based device. SR-IOV allows a physical resource (e.g., a single network interface controller) to appear as multiple resources (e.g., sixty-four network interface controllers). Thus, a PCI-based device providing a certain functionality (e.g., a network interface controller) may appear to a device making use of the PCI-based device to be multiple devices providing the same functionality. The functions of an SR-IOV-capable storage adapter device may be classified as physical functions (PFs) or virtual functions (VFs). Physical functions are fully featured functions of the device that can be discovered, managed, and manipulated. Physical functions have configuration resources that can be used to configure or control the storage adapter device. Physical functions include the same configuration address space and memory address space that a non-virtualized device would have. A

physical function may have a number of virtual functions associated with it. Virtual functions are similar to physical functions, but are light-weight functions that may generally lack configuration resources, and are generally controlled by the configuration of their underlying physical functions. Each of the physical functions and/or virtual functions may be assigned to a respective thread of execution (such as for example, a virtual machine) running on a host device.

(134) The modules described herein may be software modules, hardware modules or a suitable combination thereof. If the modules are software modules, the modules can be embodied on a non-transitory computer-readable medium and processed by a processor in any of the computer systems described herein. It should be noted that the described processes and architectures can be performed either in real-time or in an asynchronous mode prior to any user interaction. The modules may be configured in the manner suggested in the preceding figures, and/or functions described herein can be provided by one or more modules that exist as separate modules and/or module functions described herein can be spread over multiple modules.

(135) The specification and drawings are, accordingly, to be regarded in an illustrative rather than a restrictive sense. It will, however, be evident that various modifications and changes may be made thereunto without departing from the broader spirit and scope of the disclosure as set forth in the claims.

(136) Other variations are within the spirit of the present disclosure. Thus, while the disclosed techniques are susceptible to various modifications and alternative constructions, certain illustrated examples thereof are shown in the drawings and have been described above in detail. It should be understood, however, that there is no intention to limit the disclosure to the specific form or forms disclosed, but on the contrary, the intention is to cover all modifications, alternative constructions, and equivalents falling within the spirit and scope of the disclosure, as defined in the appended claims.

(137) The use of the terms “a” and “an” and “the” and similar referents in the context of describing the disclosed examples (especially in the context of the following claims) are to be construed to cover both the singular and the plural, unless otherwise indicated herein or clearly contradicted by context. The terms “comprising,” “having,” “including,” and “containing” are to be construed as open-ended terms (i.e., meaning “including, but not limited to,”) unless otherwise noted. The term “connected” is to be construed as partly or wholly contained within, attached to, or joined together, even if there is something intervening. Recitation of ranges of values herein are merely intended to serve as a shorthand method of referring individually to each separate value falling within the range, unless otherwise indicated herein and each separate value is incorporated into the specification as if it were individually recited herein. All methods described herein can be performed in any suitable order unless otherwise indicated herein or otherwise clearly contradicted by context. The use of any and all examples, or exemplary language (e.g., “such as”) provided herein, is intended merely to better illuminate examples of the disclosure and does not pose a limitation on the scope of the disclosure unless otherwise claimed. No language in the specification should be construed as indicating any non-claimed element as essential to the practice of the disclosure.

(138) Disjunctive language such as the phrase “at least one of X, Y, or Z,” unless specifically stated otherwise, is intended to be understood within the context as used in general to present that an item, term, etc., may be either X, Y, or Z, or any combination thereof (e.g., X, Y, and/or Z). Thus, such disjunctive language is not generally intended to, and should not, imply that certain examples require at least one of X, at least one of Y, or at least one of Z to each be present.

(139) Various examples of this disclosure are described herein, including the best mode known to the inventors for carrying out the disclosure. Variations of those examples may become apparent to those of ordinary skill in the art upon reading the foregoing description. The inventors expect skilled artisans to employ such variations as appropriate and the inventors intend for the disclosure to be practiced otherwise than as specifically described herein. Accordingly, this disclosure

includes all modifications and equivalents of the subject matter recited in the claims appended hereto as permitted by applicable law. Moreover, any combination of the above-described elements in all possible variations thereof is encompassed by the disclosure unless otherwise indicated herein or otherwise clearly contradicted by context.

Claims

1. A method for aligning tensors for execution of a bidirectional recurrent neural network (RNN) on a data-flow centric processor, the method comprising: compiling the bidirectional RNN for a user-specified length (L), the bidirectional RNN having a forward network, a reverse network, a first transformation layer, and a second transformation layer, wherein the forward network performs a set of operations in a forward direction and the reverse network performs the set of operations in a reverse direction; and during runtime and while performing an inference using the data-flow centric processor: determining a number of runtime steps (N), based on the inference; providing an input tensor to the bidirectional RNN network; generating, using the forward network and based on the input tensor, a forward output tensor; performing, using the first transformation layer, a first transformation on the input tensor to generate a transformed input tensor, wherein performing the first transformation using the first transformation layer includes modifying an order of the input tensor by reversing an order of only a first N values of the input tensor, the input tensor having a length greater than N; generating, using the reverse network and based on the transformed input tensor, a reverse output tensor; performing, using the second transformation layer, a second transformation on the reverse output tensor to generate a transformed output tensor, wherein performing the second transformation includes modifying an order of a first N values of the reverse output tensor; and combining the forward output tensor and the transformed output tensor to produce an output tensor.
2. The method of claim 1, wherein modifying the order of the first N values of the reverse input tensor includes: reversing the order of the first N values of the reverse output tensor.
3. The method of claim 1, wherein performing the first transformation further includes: setting the last L-N values of the input tensor to zero.
4. The method of claim 1, further comprising: during runtime and while performing the inference: determining that N is equal to or less than L.
5. A method comprising: providing an input to a bidirectional network, the bidirectional network having a forward network and a reverse network, wherein the forward network performs a set of operations in a forward direction and the reverse network performs the set of operations in a reverse direction; generating, using the forward network and based on the input, a forward output; performing, using a first transformation layer of the bidirectional network, a first transformation on the input to generate a transformed input, wherein performing the first transformation using the first transformation layer includes modifying an order of the input by reversing an order of only a first N values of the input tensor, the input tensor having a length greater than N; generating, using the reverse network and based on the transformed input, a reverse output; and performing, using a second transformation layer of the bidirectional network, a second transformation on the reverse output to generate a transformed output, wherein performing the second transformation includes modifying an order of a set of values of the reverse output.
6. The method of claim 5, further comprising: combining the forward output and the transformed output to produce an output.
7. The method of claim 5, further comprising: determining a number of runtime steps based on the input.
8. The method of claim 5, wherein modifying the order of the input includes: transposing the first N values of the input.
9. The method of claim 5, wherein modifying the order of the reverse output includes: reversing the

order of the set of values of the reverse output.

10. The method of claim 5, wherein performing the first transformation further includes: setting a last set of values of the input to zero.

11. The method of claim 5, wherein performing the second transformation further includes: overwriting a last set of values of the reverse output with zeros.

12. A non-transitory computer-readable medium having stored therein instructions that, when executed by one or more processors cause the one or more processors to perform operations including: providing an input to a bidirectional network, the bidirectional network having a forward network and a reverse network, wherein the forward network performs a set of operations in a forward direction and the reverse network performs the set of operations in a reverse direction; generating, using the forward network and based on the input, a forward output; performing, using a first transformation layer of the bidirectional network, a first transformation on the input to generate a transformed input, wherein performing the first transformation using the first transformation layer includes modifying an order of the input by reversing an order of only a first N values of the input tensor, the input tensor having a length greater than N ; generating, using the reverse network and based on the transformed input, a reverse output; and performing, using a second transformation layer of the bidirectional network, a second transformation on the reverse output to generate a transformed output, wherein performing the second transformation includes modifying an order of a set of values of the reverse output.

13. The non-transitory computer-readable medium of claim 12, wherein the operations further comprise: combining the forward output and the transformed output to produce an output.

14. The non-transitory computer-readable medium of claim 12, wherein the operations further comprise: determining a number of runtime steps based on the input.

15. The non-transitory computer-readable medium of claim 12, wherein modifying the order of the set of values of the reverse output includes: reversing the order of the set of values of the reverse output.

16. The non-transitory computer-readable medium of claim 12, wherein performing the first transformation further includes: setting a last set of values of the input to zero.

17. The non-transitory computer-readable medium of claim 12, wherein performing the second transformation further includes: overwriting a last set of values of the reverse output with zeros.
