

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent

12400645

Kind Code

B1

Date of Patent

August 26, 2025

Inventor(s)

Narayan-Chen; Anjali Yuan et al.

Virtual assistant humor management

Abstract

In accordance with one disclosed method, first data representing one or more first keywords may be processed by at least one first component to determine a first pair of words suitable for generating a pun. The first pair of words may then be processed using a first machine learning model configured to generate an output pun based at least in part on one or more input keywords and an input word pair such that the output pun is contextually related to the one or more input keywords. A first pun that is contextually related to the one or more first keywords may be received from the first machine learning model, and a first device may be caused to output a representation of the first pun.

Inventors: Narayan-Chen; Anjali Yuan (Santa Clara, CA), Sun; Jiao (Los Angeles, CA), Oraby; Shereen (Santa Clara, CA), Gao; Shuyang (Sunnyvale, CA), Chung; Tagyoung (Palo Alto, CA), Huang; Jing (Mountain View, CA), Liu; Yang (Los Altos, CA), Peng; Nanyung (Los Angeles, CA), Mandal; Arindam (Redwood City, CA), Natarajan; Premkumar (Rolling Hills Estates, CA)

Applicant: Amazon Technologies, Inc. (Seattle, WA)

Family ID: 1000006861516

Assignee: Amazon Technologies, Inc. (Seattle, WA)

Appl. No.: 18/082838

Filed: December 16, 2022

Related U.S. Application Data

us-provisional-application US 63417882 20221020

Publication Classification

Int. Cl.: G10L15/22 (20060101); G10L13/02 (20130101); G10L15/183 (20130101)

U.S. Cl.:

Field of Classification Search

CPC: G10L (15/183); G10L (15/02)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
9858925	12/2017	Gruber	N/A	N/A
10037768	12/2017	Akkiraju	N/A	N/A
10297273	12/2018	Akkiraju	N/A	N/A
10311895	12/2018	Akkiraju	N/A	N/A
10424319	12/2018	Akkiraju	N/A	N/A
10642939	12/2019	Toplyn	N/A	N/A
10878817	12/2019	Toplyn	N/A	N/A
11080012	12/2020	Lemay	N/A	N/A
11080485	12/2020	Toplyn	N/A	N/A
11562142	12/2022	Nijkamp	N/A	N/A
12211495	12/2024	Natarajan	N/A	N/A
2004/0088327	12/2003	Shimomura	N/A	G06F 40/56
2019/0096425	12/2018	Akkiraju	N/A	N/A
2019/0096426	12/2018	Akkiraju	N/A	N/A
2019/0096427	12/2018	Akkiraju	N/A	N/A
2019/0266250	12/2018	Toplyn	N/A	N/A
2019/0355381	12/2018	Akkiraju	N/A	N/A
2020/0227032	12/2019	Toplyn	N/A	N/A
2021/0103700	12/2020	Toplyn	N/A	N/A
2022/0067077	12/2021	Tupakula	N/A	G06F 16/345
2022/0277141	12/2021	Nijkamp	N/A	N/A

OTHER PUBLICATIONS

Karyn Buxman. 2008. Humor in the OR: A Stitch in Time? AORN Journal, Jul. 2008, vol. 88, No. 1, pp. 67-77. cited by applicant

Anne Cutler, et al. 1977. On the role of sentence stress in sentence processing*, Language and Speech, pp. 1-10. cited by applicant

Joseph L. Fleiss. 1973. The Equivalence of Weighted Kappa and the Intraclass Correlation Coefficient as Measures of Reliability. Educational and Psychological Measurement, 1973, vol. 33, pp. 613-619. cited by applicant

Mikyong Kim, et al. 2000. Patterns of Comprehension and Production of Nouns and Verbs in Agrammatism: Implications for Lexical Organization. Brain and Language vol. 74, pp. 1-25. Retrieved from <http://www.idealibrary.com>. cited by applicant

Jacob L. Moreno. 1955. Theory of Spontaneity-Creativity. Sociometry, vol. 18, No. 4, pp. 105-118. Retrieved from JSTOR, <https://www.jstor.org/stable/2785848>; or <https://doi.org/10.2307/2785848>. cited by applicant

Stuart Rose, et al. 2010. Automatic keyword extraction from individual documents. Text Mining: Applications and Theory, vol. 1, pp. 1-20. cited by applicant

Yufeng Diao, et al. 2019. Heterographic Pun Recognition via Pronunciation and Spelling Understanding Gated Attention Network. The World Wide Web Conference, pp. 363-371.A. cited by applicant

Abolfazl Horri. 2011. Linguistic mechanisms of humor: Pun and/or ambiguity. Language Related Research, vol. 2, Issue 2, pp. 19-40. Abstract only, retrieved from <https://scholar.google.com/citations?>

view_op=view_citation&hl=en&user=zbWiVNIAAAJ&citation_for_view=zbWiVNIAAAJ:Se3iqnhoufwC. cited by applicant

U.S. Appl. No. 18/082,790, filed Dec. 16, 2022. cited by applicant

Sura Dhiaa Ibraheem, et al. 2016. Pun and (Un)Intentional Humor. Journal of American Academic Research. Retrieved from <https://www.researchgate.net/publication/299524825>, pp. 1-18. cited by applicant

Jacob Devlin, et al. 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In Proceedings of North American Association for Computational Linguistics-Human Language Technologies 2019, pp. 4171-4186. Association for Computational Linguistics. cited by applicant

Aparna Garimella, et al. 2020. Judge me by my size (noun), do you? YodaLib: A Demographic-Aware Humor Generation Framework. In Proceedings of the 28th International Conference on Computational Linguistics (Online), pp. 2814-2825. Association for Computational Linguistics. cited by applicant

Md Kamrul Hasan, et al. 2019. UR-FUNNY: A Multimodal Language Dataset for Understanding Humor. In Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing, pp. 2046-2056. Association for Computational Linguistics. cited by applicant

Tatsunori B. Hashimoto, et al. 2019. A Retrieve-and-Edit Framework for Predicting Structured Outputs. 32nd Conference on Neural Information Processing Systems (NIPS 2018). Retrieved from <https://dl.acm.org/doi/10.5555/3327546.3327670>, pp. 1-11. cited by applicant

He He, et al. 2019. Pun Generation with Surprise. In Proceedings of NAACL HLT 2019, pp. 1734-1744. Association for Computational Linguistics. cited by applicant

Pengcheng He, et al. 2021. DeBERTa: Decoding-Enhanced BERT with Dis-Entangled Attention. Published as a conference paper at ICLR 2021. arXiv preprint, arXiv:2006.03654v6, pp. 1-23. cited by applicant

Bryan Anthony Hong, et al. 2009. Automatically Extracting Word Relationships as Templates for Pun Generation. In Proceedings of the NAACL HLT Workshop on Computational Approaches to Linguistic Creativity, pp. 24-31. Association for Computational Linguistics. cited by applicant

I-Hung Hsu, et al. 2022. Degree: A Data-Efficient Generation-Based Event Extraction Model. In Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, pp. 1890-1908. Association for Computational Linguistics. cited by applicant

Kuan-Hao Huang, et al. 2021. Generating Syntactically Controlled Paraphrases without Using Annotated Paralell Pairs. In Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics, pp. 1022-1033. Association for Computational Linguistics. cited by applicant

Mike Lewis, et al. 2020. BART: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. In Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, pp. 7871-7880. Association for Computational Linguistics. cited by applicant

Yinhan Liu, et al. 2019. RoBERTa: A Robustly Optimized BERT Pretraining Approach. arXiv preprint arXIV:1907.11692v1, 13 pages. cited by applicant

Ilya Loshchilov, et al. 2019. Decoupled Weight Decay Regularization. In 7th International Conference on Learning Representations, ICLR 2019. Retrieved from <https://openreview.net/pdf?id=Bkg6RiCqY7>, 18 pages. cited by applicant

Fuli Luo, et al. 2019. Pun-GAN: Generative Adversarial Network for Pun Generation. In Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing, pp. 3388-3393. Association for Computational Linguistics. cited by applicant

Tristan Miller, et al. 2015. Automatic disambiguation of English puns. In Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing, pp. 719-729. Association for Computational Linguistics. cited by applicant

Tristan Miller, et al. 2017. SemEval-2017 Task 7: Detection and Interpretation of English puns. In Proceedings of the 11th International Workshop on Semantic Evaluation (SemEval-2017), pp. 58-58. Association for Computational Linguistics. cited by applicant

Anirudh Mittal, et al. 2021. "So You Think You're Funny?": Rating the Humour Quotient in Standup Comedy. In Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing, pp. 10073-10079. Association for Computational Linguistics. cited by applicant

Anirudh Mittal, et al. 2022. AmbiPun: Generating humorous puns with ambiguous context. In Proceedings of

the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, pp. 1053-1062. Association for Computational Linguistics. cited by applicant

Jeffrey Pennington, et al. 2014. GloVe: Global vectors for word representation. In Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP), pp. 1532-1543. Association for Computational Linguistics. cited by applicant

Saša Petrovic, et al. 2013. Unsupervised joke generation from big data. In Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (vol. 2: Short Papers), pp. 228-232. Association for Computational Linguistics. cited by applicant

Colin Raffel, et al. 2020. Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer. Journal of Machine Learning Research, 21 (2020). arXiv preprint arXiv:1910.10683v3, 67 pages. cited by applicant

Graeme Ritchie. 2005. Computational Mechanisms for Pun Generation. In Proceedings of the Tenth European Workshop, 8 pages. Association for Computational Linguistics. cited by applicant

Jiao Sun, et al. 2021. AESOP: Paraphrase generation with adaptive syntactic control. In Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing, pp. 5176-5189. Association for Computational Linguistics. cited by applicant

Jiao Sun, et al., 2022. ExPUNations: Augmenting Puns with Keywords and Explanations. arXiv preprint arXiv:2210/13513v1, 16 pages. cited by applicant

Yufei Tian, et al. 2022. Zero-shot Sonnet Generation with Discourse-level Planning and Aesthetics Features. In Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, pp. 3587-3599. Association for Computational Linguistics. cited by applicant

Alessandro Valitutti, et al. 2013. In Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics, pp. 243-248. Association for Computational Linguistics. cited by applicant

Thomas Wolf, et al. 2020. Transformers: State-of-the-art natural language processing. In Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations, pp. 38-45, Online. Association for Computational Linguistics. cited by applicant

Ziqing Yang, et al. 2020. TextBrewer: An Open-Source Knowledge Distillation Toolkit for Natural Language Processing. In Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics: System Demonstrations, pp. 9-16, Online. Association for Computational Linguistics. cited by applicant

Zixiaofan Yang, et al. Julia Hirschberg. 2021. CHoRaL: Collecting Humor Reaction Labels from Millions of Social Media Users. In Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing, pp. 4429-4435, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics. cited by applicant

Zhiwei Yu, et al. 2018. A Neural Approach to Pun Generation. In Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (vol. 1: Long Papers), pp. 1650-1660. Association for Computational Linguistics. cited by applicant

Zhiwei Yu, et al. 2020. Homophonic Pun Generation with Lexically Constrained Rewriting. In Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing, Online, pp. 2870-2876. Association for Computational Linguistics. cited by applicant

Yukun Zhu, et al. 2015. Aligning Books and Movies: Towards Story-like Visual Explanations by Watching Movies and Reading Books, pp. 19-27. arXiv preprint arXiv:1506.06724v1, 23 pages. cited by applicant

Satanjeev Banerjee, et al. 2005. METEOR: An automatic metric for MT evaluation with improved correlation with human judgments. In Proceedings of the ACL Workshop on Intrinsic and Extrinsic Evaluation Measures for Machine Translation and/or Summarization, pp. 65-72. Association for Computational Linguistics. cited by applicant

Dana-Maria Camburu, et al. 2018. e-SNLI: Natural Language Inference with Natural Language Explanations. In Proceedings of Neural Information Processing Systems. Advances in Neural Information Processing Systems 31 (NeurIPS 2018), pp. 1-11. cited by applicant

Santiago Castro, et al. 2018. A Crowd-Annotated Spanish Corpus for Humor Analysis In Proceedings of the Sixth International Workshop on Natural Language Processing for Social Media, pp. 7-11. Association for Computational Linguistics. cited by applicant

Miruna-Adriana Clinciu, et al. 2021. A Study of Automatic Metrics for the Evaluation of Natural Language

Explanations. In Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume, pp. 2376-2387, Online. Association for Computational Linguistics. cited by applicant

Alon Jacovi, et al. 2020. Towards Faithfully Interpretable NLP Systems: How Should We Define and Evaluate Faithfulness? In Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, pp. 4198-4205. Association for Computational Linguistics. cited by applicant

Peter Jansen, et al. 2018. WorldTree: A Corpus of Explanation Graphs for Elementary Science Questions supporting Multi-Hop Inference. In Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018), pp. 2732-2740. European Language Resources Association (ELRA). cited by applicant

Justine T. Kao, et al. 2016. A Computational Model of Linguistic Humor in Puns. Cognitive Science, 40(5), pp. 1270-1285. Retrieved from <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC5042108>. cited by applicant

Maxime Kayser, et al. 2021. e-ViL: A Dataset and Benchmark for Natural Language Explanations in Vision-Language Tasks. e-ViL: A Dataset and Benchmark for Natural Language Explanations in Vision-Language Tasks, pp. 1-24. arXiv preprint arXiv:2105.03761v2. cited by applicant

Sawan Kumar, et al. 2020. NILE : Natural Language Inference with Faithful Natural Language Explanations. In Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, pp. 8730-8742. Association for Computational Linguistics. cited by applicant

Wang Ling, et al. 2017. Program Induction by Rationale Generation: Learning to Solve and Explain Algebraic Word Problems. In Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics, pp. 158-167. Association for Computational Linguistics. cited by applicant

Binny Mathew, et al. 2022. HateXplain: A Benchmark Dataset for Explainable Hate Speech Detection*. arXiv preprint arXiv 2012.10289v.2, 12 pages. cited by applicant

Kishore Papineni, et al. 2002. BLEU: a Method for Automatic Evaluation of Machine Translation. In Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics, pp. 311-318. Association for Computational Linguistics. cited by applicant

Nazneen Fatema Rajani, et al. 2019. Explain yourself! Leveraging Language Models for Commonsense Reasoning. In Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics, pp. 4932-4942. Association for Computational Linguistics. cited by applicant

Jiao Sun, et al. 2022. Context-Situated Pun Generation. arXiv preprint arXiv:2210.13522v1, 14 pages. cited by applicant

Orion Weller, et al. 2019. Humor Detection: A Transformer Gets the Last Laugh. In Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing, pp. 3621-3625. Association for Computational Linguistics. cited by applicant

Sarah Wiegrefe, et al. 2021. Measuring Association Between Labels and Free-Text Rationales. In the Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing, 19 pages. Retrieved from <https://aclanthology.org/2021.emnlp-main.804.pdf>. cited by applicant

Wangchunshu Zhou, et al. 2022. Towards Interpretable Natural Language Understanding with Explanations as Latent Variables. 34th Conference on Neural Information Processing Systems (NeurIPS 2020), pp. 1-16. arXiv preprint arXiv:2011.05268v3. cited by applicant

Yanyan Zou, et al. 2019. Joint Detection and Location of English Puns. In Proceedings of NAACL-HLT 2019, pp. 2117-2123. Association for Computational Linguistics. cited by applicant

Office Action mailed Feb. 26, 2025 for U.S. Appl. No. 18/082,790. cited by applicant

Final Office Action issued Jun. 6, 2025 for U.S. Appl. No. 18/082,790. cited by applicant

Primary Examiner: Siddo; Ibrahim

Attorney, Agent or Firm: Pierce Atwood LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION (1) This application claims the benefit of and priority to U.S. Provisional Patent Application No. 63/417,882, filed Oct. 20, 2022, and entitled “VIRTUAL ASSISTANT PUN GENERATION AND MANAGEMENT,” in the names of Anjali Yuan Narayan-Chen, et al. The above provisional application is herein incorporated by reference in its entirety.

BACKGROUND

(1) Natural language processing systems have progressed to the point where humans can interact with and control computing devices using their voices. Such systems employ techniques to identify the words spoken by a user based on the various qualities of received input data. Speech recognition combined with natural language understanding processing techniques enable speech-based user control of computing devices to perform tasks based on the spoken inputs. Speech recognition and natural language understanding processing techniques are sometimes referred to collectively or separately as spoken language understanding (SLU) processing. SLU processing may be used by computers, hand-held devices, telephone computer systems, kiosks, and a wide variety of other devices to improve human-computer interactions.

Description

BRIEF DESCRIPTION OF DRAWINGS

- (1) For a more complete understanding of the present disclosure, reference is now made to the following description taken in conjunction with the accompanying drawings.
- (2) FIG. 1A illustrates an example system configured to recognize whether and/or why content is humorous in accordance with some embodiments of the present disclosure;
- (3) FIG. 1B illustrates an example system configured to generate humorous content based on keywords representing a current context, such as content of an ongoing dialog between a user and a voice-controlled electronic device, in accordance with some embodiments of the present disclosure;
- (4) FIG. 2A is a block diagram illustrating example components that may be included in the humor recognizer shown in FIG. 1A;
- (5) FIG. 2B is a table showing example data that may be used to train the pun explanation model and/or the pun classification model shown in FIG. 2A;
- (6) FIG. 2C illustrates an example process for training the pun explanation model shown in FIG. 2A using training data from the table shown in FIG. 2B;
- (7) FIG. 2D illustrates an example process for training the pun classification model shown in FIG. 2A using training data from the table shown in FIG. 2B;
- (8) FIG. 3A is a block diagram illustrating example components that may be included in the humor generator shown in FIG. 1B;
- (9) FIG. 3B is a table showing example data that may be used to train the pun word classification model and/or the pun generation model shown in FIG. 3A;
- (10) FIG. 3C illustrates an example process for training the pun word classification model shown in FIG. 3A using training data from the table shown in FIG. 3B;
- (11) FIG. 3D illustrates an example process for training the pun generation model shown in FIG. 3A using training data from the table shown in FIG. 3B;
- (12) FIG. 4 is a conceptual diagram of components of the system shown in FIGS. 1A and 1B, according to some embodiments of the present disclosure;
- (13) FIG. 5 is a conceptual diagram illustrating components that may be included in the device shown in FIGS. 1A and 1B, according to some embodiments of the present disclosure;
- (14) FIG. 6 shows an example implementation of a dialog management component, such as the dialog management components shown in FIGS. 4 and 5, according to embodiments of the present disclosure;
- (15) FIG. 7 shows an example process that may be performed by one or more components of the dialog management component shown in FIG. 6 to recognize and respond to humor during a dialog;
- (16) FIG. 8 shows an example process that may be performed by one or more components of the dialog management component shown in FIG. 6 to inject context-based humor into a dialog at an appropriate

time;

(17) FIG. 9 is a block diagram conceptually illustrating example components of a device, according to some embodiments of the present disclosure;

(18) FIG. 10 is a block diagram conceptually illustrating example components of a system, according to some embodiments of the present disclosure; and

(19) FIG. 11 illustrates an example of a computer network for use with the overall system, according to some embodiments of the present disclosure.

DETAILED DESCRIPTION

(20) Automatic speech recognition (ASR) is a field of computer science, artificial intelligence, and linguistics concerned with transforming audio data associated with speech into a token or textual representation of that speech. Natural language understanding (NLU) is a field of computer science, artificial intelligence, and linguistics concerned with enabling computers to derive meaning from natural language inputs (such as spoken inputs). ASR and NLU are often referred to collectively as spoken language understanding (SLU). Text-to-speech (TTS) is a field of computer science concerning transforming textual and/or other data into audio data that is synthesized to resemble human speech.

Natural language generation (NLG) is a field of computer science concerning generation of text from structured data, where the text represents meaningful phrases and sentences in a natural language form.

(21) Humor serves multiple purposes and provides numerous benefits, such as relieving anxiety, avoiding painful feelings, and facilitating learning. As a specific example of humor, the creative uses of puns, wordplay, and ambiguity are important ways to come up with jokes. Pun understanding and generation are particularly challenging tasks to implement for a virtual assistant/chatbot system because they require extensive commonsense and world knowledge to compose and understand, even for humans. Despite growing interest in the area, there are limited amounts of data available in the domain of humor understanding and generation. Existing humor datasets are typically annotated only with binary labels indicating whether each sentence is a joke, pun, or punchline.

(22) Such binary labels may be insufficient to benchmark the ability of machine learning models (“models”) or other components to understand and generate humorous text, as such sparse supervision signals and coarse-grained annotations may provide insufficient information to perform system configuration/training to effectively incorporate humor aspects in a system-to-user interaction. To account for these shortcomings, offered are techniques for the formulation of a significantly more robust humor dataset that can be used to effectively train (and operate at runtime) models to recognize whether and/or why text is funny. Also offered are techniques to generate humorous text based on keywords representing a current context, such as content of an ongoing dialog between a user and a voice-controlled electronic device.

Such tools may be incorporated into computing systems to both detect humorous text and to generate humorous text, enabling the system to interact with humans in a more realistic manner. For example, as outlined below, virtual assistant systems, chatbots, etc., that are adapted to engage in a dialog with a user may be configured to recognize appropriate times to inject humor into the dialog, and then generate humorous content that makes sense within the context of the dialog. Accordingly, by employing the techniques described herein, such systems may effectively be instilled with a human-like “sense of humor,” thus making the experience of interacting with such system more natural and enjoyable for users.

(23) A system according to the present disclosure may be configured to incorporate user permissions and may only perform activities disclosed herein if approved by a user. As such, the systems, devices, components, and techniques described herein would be typically configured to restrict processing where appropriate and only process user data in a manner that ensures compliance with all appropriate laws, regulations, standards, and the like. The system and techniques can be implemented on a geographic basis to ensure compliance with laws in various jurisdictions and entities in which the components of the system and/or user are located.

(24) FIGS. 1A and 1B show an example system **100** configured to recognize and/or generate humorous text. In particular, as illustrated, the system **100** may include, among other components, a humor management component **130** that includes a humor recognizer **140** and/or a humor generator **150**. FIG. 1A shows an example process **160** that may be performed by the humor recognizer **140**, and FIG. 1B shows an example process **180** that may be performed by the humor generator **150**. Those processes **160**, **180** are both described in more detail below.

(25) As illustrated in FIGS. 1A and 1B, in some implementations, the system **100** may further include a device **110** (e.g., local to a user **5**) and a system **120** (e.g., remote from the user **5**), with the device **110** being in communication with the system **120** across one or more networks **199**. The network(s) **199** may include the Internet and/or any other wide- or local-area network, and may include wired, wireless, and/or cellular network hardware. While the user **5** is illustrated as a human, it should be appreciated that the present disclosure is not limited thereto, and that the user **5** may be a non-human, such as an application/skill, bot, or the like. Further, it should be understood that although the illustrated example shows the humor management component **130** as included within the system **120**, some or all of the functionality of the humor management component **130** may additionally or alternatively be implemented elsewhere in the system **100**, such as within the device **110**.

(26) The device **110** may receive audio of a spoken natural language input from the user **5**. The device **110** may generate audio data corresponding to the audio, and may send that audio data to the system **120**. The system **120** may receive the audio data corresponding to the spoken natural language input, and may perform ASR and NLU processing on the audio data to determine an appropriate response. In some circumstances, the system **120** may generate responsive text data and perform text-to-speech processing on that text data to generate audio data. The system **120** may then send the generated audio data to the device **110** to cause the device **110** to output an audio response to the user **5**. Interactions between the device **110** and system **120** may also be graphical and/or textual based, such as the user inputting text into a graphical user interface of the device **110** and the device **110** sending the text to the system **120** and receiving output text (or other output data, such as display data) in return.

(27) In connection with such exchanges, situations can arise in which the user **5** says something funny, or in which content accessed and/or generated by the system **120**, for output by the device **110** or otherwise, includes humorous text. As described in more detail below, in some implementations, one or more models of the humor recognizer **140** can be trained and used to recognize whether and/or why text corresponding to such user input/accessed content is funny. Examples of models that may be employed for such purposes, as well as training data that can be used to train those models, are described below in connection with FIGS. 2A-D. In some implementations, in response to the humor recognizer **140** recognizing funny text, one or more suitable actions can be taken by the system **120**. For example, if the user **5** utters a sentence that itself includes a pun, or if, in response to output audio, the user **5** utters a response that includes a play on one or more of the words of the output audio, the system **120** may cause the device **110** to output “I see what you did there!,” or a similar phrase, to output a laughing or giggling sound, and/or to output a joke relating to some aspect of the detected humor. In some implementations, the humor generator **150** may monitor text (or other data corresponding to textual representation such as tokens corresponding to text) representing an utterance by the user **5**, a response provided by the system **120** to the device **110**, multiple turns of an ongoing dialog with the user **5**, content accessed by the system **120**, etc., to determine one or more or more keywords corresponding to such text. Such data may be in the form of dialog data, for example as determined by dialog management component **485** discussed below. In some implementations, a keyword extraction algorithm, such as the rapid automatic keyword extraction (RAKE) algorithm, may be used to extract one or more keywords from one or more sentences of text, such as those discussed above. In other implementations, the user **5** may utter a request to the device **110**, possibly preceded by a wakeword (e.g., “Alexa”), asking the system **100** to output a pun relating to a particular topic or that relates to one or more particular words (e.g., “Alexa, tell me a pun about cats and dogs”), and the system **100** may parse at least a portion of the utterance (e.g., a “command” portion that follows the wakeword) to identify one or more pertinent keywords to use (e.g., “cats” and “dogs”). In any such case, the determined keyword(s) may then be processed by one or more trained models to generate text representing a pun that makes sense in the context of such keyword(s). Examples of models that may be employed to generate puns based on one or more keywords, as well as data that can be used to train those models, are described below in connection with FIGS. 3A-D.

(28) As noted above, in some implementations, one or more models of the humor recognizer **140** (see FIG. 1A) can be trained and used to recognize whether and/or why text corresponding to user input/accessed content is funny. FIG. 2A illustrates two example models (e.g., a pun explanation model **202** and a pun classification model **204**) that may be separately trained and then operated together to achieve such objectives. FIG. 2B shows a table **206** including a sampling of an example training data set

that may be used to train the pun explanation model **202** and/or the pun classification model **204** in accordance with some embodiments. As shown, the table **206** may include text entries **208**, pun indicator entries **210**, and explanation entries **212**.

(29) The text entries **208** in the table **206** may each include the text of one or more sentences, and the pun indicator entries **210** in the table **206** may indicate whether the corresponding text entries **208** represent puns. For example, in some implementations, a pun indicator entry **210** with Boolean value of “1” may indicate that the corresponding text entry **208** (e.g., “T1” which may correspond to a series of words) represents a pun, whereas a pun indicator entry **210** with Boolean value of “0” may indicate that the corresponding text entry **208** (e.g., “T2” which may correspond to a different series of words) does not represent a pun. The explanation entries **212** may represent human generated explanations for those text entries **208** that represent puns, e.g., the text entries **208** for which the pun indicator entry **210** includes a Boolean value of “1.” As one example, the text entry “T1” in the table **206** could be the text “When artists dream in color it's a pigment of their imagination,” and the explanation “E1” in the table **206** could be the text “Pigments are non-soluble materials often used in painting, and pigment sounds like figment, which is something that is not real but some believes it is.” As another example, the text entry “T3” in the table **206** could be the text “The man found something to catch fish, which was a net gain,” and the explanation “E3” in the table **206** could be the text “This is a play on words. A ‘net gain’ means an increase in revenue but here ‘net’ refers to how a net is used to catch fish.” For text entries **208** that do not represent puns (e.g., “T3” in the table **206**), the explanation entries **212** may include the same values as the corresponding text entries **208**. Illustrative techniques for accumulating and populating the table **206** with a sufficiently large quantity of suitable data to train the pun explanation model **202** and/or the pun classification model **204** is described in more detail in Appendix A of U.S. Provisional Patent Application No. 63/417,882.

(30) As illustrated in FIG. 2A, after the models **202** and **204** have been trained, as described below in connection with FIGS. 2C and 2D, respectively, the pun explanation model **202** may receive input text (e.g., one or more sentences), together with added prefix text (e.g., “Generate an explanation of the pun:”) indicating the operation the pun explanation model **202** is to perform (e.g., based on the pun explanation model **202** having been fine-tuned for generating pun explanations using that same prefix text, as described below), and the pun explanation model **202** may process the input text (e.g., as a set of tokens representing words) to generate a natural language explanation of the pun, such as a description of a reason why the pun is funny and/or interesting. Further, as also illustrated in FIG. 2A, in some implementations, the generated pun explanation may then be input to the pun classification model **204**, together with the input text that was provided to the pun explanation model **202**, and the pun classification model **204** may process those inputs (e.g., as a set of tokens representing words of the input text and the pun explanation) to generate an indication (e.g., a Boolean output) of whether the input text represents a pun or does not represent a pun.

(31) FIG. 2C illustrates an example process for training the pun explanation model **202** shown in FIG. 2A. In some implementations, the pun explanation model **202** may comprise a text-to-text transfer transformer (or “T5”) model that has been pre-trained on a large corpus of textual information and fine-tuned using certain data from the table **206** (shown in FIG. 2B) and an appropriate text prefix (e.g., “Generate an explanation of the pun:”). Various techniques for pre-training and fine-tuning a T5 model are known in the art. As shown in FIG. 2C, a training module **214** may be used to train (or fine-tune) the pun explanation model **202** (e.g., a pre-trained T5 model) by identifying the text entries **208** in the table **206** that represent puns (e.g., based on the values of the pun indicator entries **210**), and using those text entries, together with a suitable prefix (e.g., “Generate an explanation of the pun:”) and the corresponding explanation entries **212** in the table **206** to train/fine-tune the pun explanation model **202**, with the text of the explanation entries **212** being used to supervise the training process.

(32) FIG. 2D shows an example process for training the pun classification model **204** shown in FIG. 2A. In some implementations, the pun classification model **204** may be a bidirectional encoder representations from transformers (or “BERT”) model that has been pre-trained and fine-tuned using data from the table **206** (shown in FIG. 2B). Various techniques for pre-training and fine-tuning a BERT model are known in the art. For example, as FIG. 2D illustrates, in some implementations, a training module **216** may be used to train or fine-tune the pun classification model **204**, e.g., a pre-trained BERT

model, using the values in respective rows of table **206**, with the values of the pun indicator entries **210** being used to supervise the training process.

(33) As noted above, in some implementations, the humor recognizer **140** may perform the example process **160** (shown in FIG. 1A) to determine whether first text corresponds to a pun and to cause the device **110** to take at least one action, e.g., to output a laugh, a groan, a response recognizing the pun (e.g., “I see what you did there”), another pun contextually related to the detected pun, etc.

(34) At a step **162** of the process **160**, the humor management component **130** may receive first data representing the first text. The first text can come from any of a number of sources. In some implementations, for example, at least a portion of the first data can be generated based on audio data representing an utterance by the user **5**.

(35) Pursuant to steps **164** and **166** of the process **160**, a first machine learning model of the humor recognizer **140** may process the first data, and output second data representing an explanation as to why the first text is humorous. In some implementations, for example, the pun explanation model **202** described in connection with FIGS. 2A and 2C may be used as the first machine learning model.

(36) Pursuant to steps **168** and **170** of the process **160**, a second machine learning model of the humor recognizer **140** may process the first data and the second data, and output a first value indicating that the first data and the second data correspond to a pun. In some implementations, for example, the pun classification model **204** described in connection with FIGS. 2A and 2D may be used as the second machine learning model.

(37) At the step **172** of the process **160**, the humor recognizer **140** may cause a first device (e.g., the device **110**) to take at least a first action based at least in part on the first value output by the second machine learning model. In some implementations, for example, the humor recognizer **140** may prompt a component of a dialog management component **485/585** (e.g., the response management component **620** shown in FIG. 6) to generate audio data corresponding to an acknowledgement of the detected pun, e.g., a laugh, a groan, etc., and may send that audio data to the device **110** to cause the device **110** to output corresponding audio.

(38) As noted above, one or more models of the humor generator **150** (see FIG. 1B) may be trained to process one or more keywords to generate text representing a pun that makes sense in the context of such keyword(s). FIG. 3A illustrates two example models (e.g., a pun word classification model **304** and a pun generation model **308**) that may be separately trained and then operated together to achieve that objective. As shown, in some implementations, a keyword extractor **302** may additionally be employed ahead of the pun word classification model **304**, and a pun word selector **306** may additionally be deployed between the pun word classification model **304** and the pun generation model **308**.

(39) FIG. 3B shows a table **310** including a sampling of an example training data set that may be used to train the pun word classification model **304** and/or the pun generation model **308** in accordance with some embodiments. As shown, in some implementations, the table **310** may include keyword(s) entries **312**, pun word entries **314**, pun word sense entries **316** (e.g., describing meanings or definitions of the pun words represented by the corresponding pun word entries **314**), alternate word entries **318**, alternate word sense entries **320** (e.g., describing meanings or definitions of the alternate words represented by the corresponding alternate word entries **318**), pun suitability entries **322**, and pun entries **324**. For each row of the table **310**, the pun word represented by the pun word entry **314** (e.g., PW1) and the alternate word represented by the alternative word entry **318** (e.g., AW1) may together form a “pun word pair” that is associated with the one or more keywords represented by the keyword(s) entry **312** (e.g., K1, K2) in that same row. The pun word sense entry **316** (e.g., PWS1) may represent a meaning or definition of the pun word represented by the corresponding pun word entry **316** (e.g., PW1), and the alternate word sense entry **320** (e.g., AWS1) may represent a meaning or definition of the alternate word represented by the corresponding alternate word entry **318** (e.g., AW1). As an illustrative example, K1 and K2 may be “construction” and “workers,” respectively, PW1 may be “stair,” PWS1 may be “support consisting of a place to rest the foot while ascending or descending a stairway,” AW1 may be “stare,” AWS1 may be “look at with fixed eyes,” and P1 may be “Two construction workers had a staring contest.” The pun suitability indicator entries **322** may indicate whether the table **310** includes suitable puns (e.g., P1) corresponding to the pun word pairs (e.g., PW1 and AW1) and the keyword(s) (e.g., K1 and K2) in the respective rows.

(40) As illustrated in FIG. 3A, after the models **304** and **308** have been trained, as described below in connection with FIGS. 3C and 3D, respectively, the pun classification model **304** may receive one or more keywords from the keyword extractor **302**, and may process the received keyword(s) to determine one or more pun word pairs that are compatible with such keyword(s). For each pun word pair that is identified as compatible, the pun word classification model **304** may output a confidence value indicating a level of confidence of that the pun word pair is compatible with the keyword(s). The pun word selector **306** may then select one or more pun word pairs having the highest confidence level(s), and may provide those pun word pairs and (optionally) their respective senses to the pun generation model **308**. The pun generation model **308** may then process those pun word pairs (and, optionally, senses), together with the keyword(s) received from the keyword extractor **302** as well as an added text prefix indicating the operation the pun generation model **308** is to perform (e.g., “Generate a pun using . . . in context . . .”), to generate a suitable pun. Similar to the pun explanation model **202** described above, the pun generation model **308** may have been fine-tuned for generating puns using that same prefix text (e.g., “Generate a pun using . . . in context . . .”).

(41) FIG. 3C shows an example process for training the pun word classification model **304** shown in FIG. 3A. In some implementations, similar to the pun classification model **204** described above, the pun word classification model **304** may be a BERT model that has been pre-trained and fine-tuned using data from the table **310** (shown in FIG. 3B). As noted above, various techniques for pre-training and fine-tuning a BERT model are known in the art. For example, as FIG. 3C illustrates, in some implementations, a training module **326** may be used to train or fine-tune the pun word classification model **304**, e.g., a pre-trained BERT model, using the keyword(s) and corresponding pun word pairs and pun suitability indicators from respective rows of the table **310**, with the pun word pairs and pun suitability indicators being used to supervise the training process. Alternatively, the pun word classification model **304** may use an unsupervised method, e.g., using a natural language inferencing (NLI) model (which may be part of an NLU component such as **460** or part of some other component), to quantify the semantic similarity between the keyword(s) and the respective pun word pairs. In either case, the pun word classification model **304** may identify pun word pairs meeting a threshold level of suitability for generating a pun in the context of the keyword(s), and may additionally or alternatively output an indicator of a degree of confidence that those pun word pairs are suitable for generating a pun in the context of the keyword(s). Pun word pairs for which the degree of confidence determined by the pun word classification model **304** exceeds the threshold may thus be deemed suitable for generating a pun in the context of the keyword(s), whereas pun word pairs for which the degree of confidence determined by the pun word classification model **304** does not exceed the threshold may be deemed unsuitable for generating a pun in the context of the keyword(s).

(42) FIG. 3D illustrates an example process for training the pun generation model **308** shown in FIG. 3A. In some implementations, similar to the pun explanation model **202** described above, the pun generation model **308** may comprise a T5 model that has been pre-trained and fine-tuned using certain data from the table **310** (shown in FIG. 3B) and an appropriate text prefix. As noted above, various techniques for pre-training and fine-tuning a T5 model are known in the art. As shown in FIG. 3D, a training module **328** may be used to train (or fine-tune) the pun generation model **308** (e.g., a pre-trained T5 model) by identifying the rows in the table **310** that include puns (e.g., based on the values of the pun suitability indicator entries **322**), and using the entries in those rows, together with a suitable prefix (e.g., “Generate a pun using . . . in context . . .”), to train or fine-tune the pun generation model **308**, with the text of the pun entries **324** being used to supervise the training process. In some implementations, a T5 model may additionally be pretrained using a corpus of non-pun sentences that use a pun word of a pun word pair. An example technique for performing such additional pre-training of a T5 model is described in Appendix B of U.S. Provisional Patent Application No. 63/417,882.

(43) In other implementations, the pun generation model **308** may employ an AmbiPun model of the type described in Mittal, et al., “AmbiPun: Generating Puns with Ambiguous Context,” 2022 *Annual Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, arXIV:2205.01825v1 (2022).

(44) As noted above, in some implementations, the humor generator **150** may perform the example process **180** (shown in FIG. 1B) to generate a pun that is contextually related to one of more particular

keywords, and to cause the device **110** to output a representation of the generated pun. As explained in more detail below, in some implementations, the generation of a pun pursuant to the process **180** may be triggered based on a determination, based on contextual data from one or more sources, that it is a good time to cause the device **110** to output a representation of a pun.

(45) At a step **182** of the process **180**, the humor generator **150** may determine first data representing one or more first keywords. In some implementations, for example, one or more sentences of a dialog between the user **5** and the system **100** may be monitored for keywords that can be used to identify suitable pun word pairs for use in generating puns that are contextually related to the keywords.

(46) At a step **184** of the process **180**, a first machine learning model may process the first data to identify a first pair of words (e.g., a pun word pair) suitable for generating a pun. In some implementations, for example, the pun word classification model **304** shown in FIGS. **3A** and **3C** may be used as the first machine learning model of the step **184**, and the pun word selector **306** may process the output of the pun word classification model **304** to select an appropriate pun word pair for use in generating a pun.

(47) At steps **186** and **188** of the process **180**, a second machine learning model may process the first pair of words (e.g., the pun word pair selected by the pun word selector **306**) to generate a first pun that is contextually related to the one or more first keywords. In some implementations, for example, the pun generation model **308** shown in FIGS. **3A** and **3D** may be used to generate the first pun.

(48) At a step **190** of the process **180**, the humor generator **150** may cause a first device, e.g., the device **110**, to output a representation of the first pun. In some implementations, for example, the humor generator **150** may send text data representing the first pun to a component of a dialog management component **485/585** (e.g., the response management component **620** shown in FIG. **6**) and that component may perform speech synthesis (e.g., TTS processing) on the text data (or instruct another speech synthesis component, e.g., a TTS component **480/580**, to perform TTS processing on the text data) to generate audio data corresponding to the first pun. Such a dialog manager may then send the generated audio data to the device **110** to cause the device **110** to output corresponding audio. FIG. **4** shows example components that may be included in the system **100** shown in FIGS. **1A** and **1B** in accordance with some embodiments. The various components illustrated in FIG. **4** may be located on the same physical device or on different physical devices. Communication between various components may occur directly or across one or more network(s) **199**.

(49) As shown in FIG. **4**, a microphone or array of microphones (of or otherwise associated with the device **110**) may capture audio **401**. The device **110** may process audio data, representing the audio **401**, to determine whether speech is detected. The device **110** may use various techniques to determine whether audio data includes speech. In some implementations, for example, the device **110** may apply voice activity detection (VAD) techniques. Such techniques may determine whether speech is present in audio data based on various quantitative aspects of the audio data, such as the spectral slope between one or more frames of the audio data, the energy levels of the audio data in one or more spectral bands, the signal-to-noise ratios of the audio data in one or more spectral bands, or other quantitative aspects. In other implementations, the device **110** may additionally or alternatively implement a classifier configured to distinguish speech from background noise. The classifier may be implemented by techniques such as linear classifiers, support vector machines, and decision trees. In still other implementations, the device **110** may additionally or alternatively apply Hidden Markov Model (HMM) or Gaussian Mixture Model (GMM) techniques to compare the audio data to one or more acoustic models in storage, which acoustic models may include models corresponding to speech, noise (e.g., environmental noise or background noise), or silence. Still other techniques may be used to determine whether speech is present in audio data.

(50) Once speech is detected in audio data, the device **110** may determine if the speech is directed at the device **110**/system **120**. In at least some embodiments, such determination may be made using a wakeword detection component **420** of the device **110**. The wakeword detection component **420** may be configured to detect various wakewords. In at least some examples, each wakeword may correspond to a name of a different digital assistant. An example wakeword/digital assistant name is “Alexa.”

(51) Wakeword detection may be performed without performing linguistic analysis, textual analysis, or semantic analysis. Instead, the audio data may be analyzed to determine if specific characteristics of the audio data match preconfigured acoustic waveforms, audio signatures, or other data corresponding to a

wakeword.

(52) Thus, the wakeword detection component **420** may compare audio data to stored data to detect a wakeword. One approach for wakeword detection applies general large vocabulary continuous speech recognition (LVCSR) systems to decode audio signals, with wakeword searching being conducted in the resulting lattices or confusion networks. Another approach for wakeword detection builds HMMs for each wakeword and non-wakeword speech signals, respectively. The non-wakeword speech includes other spoken words, background noise, etc. There can be one or more HMMs built to model the non-wakeword speech characteristics, which are named filler models. Viterbi decoding is used to search the best path in the decoding graph, and the decoding output is further processed to make the decision on wakeword presence. This approach can be extended to include discriminative information by incorporating a hybrid deep neural network (DNN)-HMM decoding framework. In another example, the wakeword detection component **420** may be built on DNN/recursive neural network (RNN) structures directly, without HMM being involved. Such an architecture may estimate the posteriors of wakewords with context data, either by stacking frames within a context window for DNN, or using RNN. Follow-on posterior threshold tuning or smoothing is applied for decision making. Other techniques for wakeword detection, such as those known in the art, may also be used.

(53) Once the wakeword detection component **420** detects a wakeword, the device **110** may “wake” and begin transmitting audio data **411**, representing the audio **401**, to the system **120**. The audio data **411** may include the detected wakeword, or the device **110** may remove the portion of the audio data **411** corresponding to the detected wakeword prior to sending the audio data **411** to the system **120**.

(54) The system **120** may include an orchestrator component **430** configured to, among other things, coordinate data transmissions between components of the system **120**. The orchestrator component **430** may receive the audio data **411** from the device **110**, and may send the audio data **411** to an ASR component **450**.

(55) The ASR component **450** may transcribe the audio data **411** into ASR output data including one or more ASR hypotheses. An ASR hypothesis may be configured as a textual interpretation of the speech in the audio data **411**, or may be configured in another manner, such as one or more tokens. Each ASR hypothesis may represent a different likely interpretation of the speech represented in the audio data **411**. Each ASR hypothesis may be associated with a score (e.g., confidence score, probability score, or the like) representing the associated ASR hypothesis correctly represents the speech in the audio **401**.

(56) The ASR component **450** may interpret the speech in the audio data **411** based on a similarity between the audio data **411** and pre-established language models. For example, the ASR component **450** may compare the audio data **411** with models for sounds (e.g., subword units, such as phonemes, etc.) and sequences of sounds to identify words that match the sequence of sounds of the speech represented in the audio data **411**.

(57) In at least some instances, instead of the device **110** receiving a spoken natural language input, the device **110** may receive a textual (e.g., typed using a keyboard) natural language input. The device **110** may determine text data representing the textual natural language input, and may send the text data to the system **120**, wherein the text data may be received by the orchestrator component **430**. The orchestrator component **430** may send the text data or ASR output data, depending on the type of natural language input received, to a NLU component **460**.

(58) The NLU component **460** may process the ASR output data or text data to determine one or more NLU hypotheses embodied in NLU output data. The NLU component **460** may perform intent classification (IC) processing on the ASR output data or text data to determine an intent of the natural language input. An intent may correspond to an action to be performed that is responsive to the natural language input. To perform IC processing, the NLU component **460** may communicate with a database of words linked to intents. For example, a music intent database may link words and phrases such as “quiet,” “volume off,” and “mute” to a <Mute> intent. The NLU component **460** may identify intents by comparing words and phrases in ASR output data or text data to the words and phrases in an intents database. In some embodiments, the NLU component **460** may communicate with multiple intents databases, with each intents database corresponding to one or more intents associated with a particular skill. A “skill” may refer to software, that may be placed on a machine or a virtual machine (e.g., software that may be launched in a virtual instance when called), configured to process NLU output data and

perform one or more actions in response thereto.

(59) As one example, IC processing of the natural language input “play my workout playlist” may determine an intent of <PlayMusic>. As another example, IC processing of the natural language input “call mom” may determine an intent of <Call>. As yet another example, IC processing of the natural language input “call mom using video” may determine an intent of <VideoCall>. In still another example, IC processing of the natural language input “what is today's weather” may determine an intent of <OutputWeather>.

(60) The NLU component **460** may also perform named entity recognition (NER) processing on the ASR output data or text data to determine one or more portions, sometimes referred to as slots, of the natural language input that may be needed for post-NLU processing (e.g., processing performed by a skill). As one example, named entity recognition (NER) processing of the natural language input “play [song name]” may determine an entity type of “SongName” and an entity value corresponding to the indicated song name. As another example, NER processing of the natural language input “call mom” may determine an entity type of “Recipient” and an entity value corresponding to “mom.” As still another example, NER processing of the natural language input “what is today's weather” may determine an entity type of “Date” and an entity value of “today.”

(61) In some embodiments, the intents identifiable by the NLU component **460** may be linked to one or more grammar frameworks with entity types that can be populated with entity values. Each entity type of a grammar framework may correspond to a portion of ASR output data or text data that the NLU component **460** identified as corresponding to an entity value. For example, a grammar framework corresponding to a <PlayMusic> intent may correspond to sentence structures such as “Play {Artist Name},” “Play {Album Name},” “Play {Song name},” “Play {Song name} by {Artist Name},” etc.

(62) For example, the NLU component **460** may perform NER processing to identify words in ASR output data or text data as subject, object, verb, preposition, etc., based on grammar rules and/or models. Then, the NLU component **460** may perform IC processing using the identified verb to identify an intent. Thereafter, the NLU component **460** may again perform NER processing to determine a grammar model associated with the identified intent. For example, a grammar model for a <PlayMusic> intent may specify a list of entity types applicable to play the identified “object” and any object modifier (e.g., a prepositional phrase), such as {Artist Name}, {Album Name}, {Song name}, etc. The NER processing may then involve searching corresponding fields in a lexicon, attempting to match words and phrases in the ASR output data that NER processing previously tagged as a grammatical object or object modifier with those identified in the lexicon.

(63) NER processing may include semantic tagging, which is the labeling of a word or combination of words according to their type/semantic meaning. NER processing may include parsing ASR output data or text data using heuristic grammar rules, or a model may be constructed using techniques such as hidden Markov models, maximum entropy models, log linear models, conditional random fields (CRFs), and the like. For example, NER processing with respect to a music skill may include parsing and tagging ASR output data or text data corresponding to “play mother's little helper by the rolling stones” as {Verb}: “Play,” {Object}: “mother's little helper,” {Object Preposition}: “by,” and {Object Modifier}: “the rolling stones.” The NER processing may identify “Play” as a verb based on a word database associated with the music skill, which IC processing determines corresponds to a <PlayMusic> intent.

(64) The NLU component **460** may generate NLU output data including one or more NLU hypotheses, with each NLU hypothesis including an intent and optionally one or more entity types and corresponding entity values. In some embodiments, the NLU component **460** may perform IC processing and NER processing with respect to different skills. One skill may support the same or different intents than another skill. Thus, the NLU output data may include multiple NLU hypotheses, with each NLU hypothesis corresponding to IC processing and NER processing performed on the ASR output or text data with respect to a different skill.

(65) As described above, in some implementations, the system **120** may perform speech processing using two different components (e.g., the ASR component **450** and the NLU component **460**). In other implementations, the system **120** may additionally or alternatively implement a spoken language understanding (SLU) component **440** configured to process audio data **411** to determine NLU output data.

(66) The SLU component **440** may be equivalent to a combination of the ASR component **450** and the NLU component **460**. Yet, the SLU component **440** may process audio data **411** and directly determine the NLU output data, without an intermediate step of generating ASR output data. As such, the SLU component **440** may take audio data **411** representing a spoken natural language input and attempt to make a semantic interpretation of the spoken natural language input. That is, the SLU component **440** may determine a meaning associated with the spoken natural language input and then implement that meaning. For example, the SLU component **440** may interpret audio data **411** representing a spoken natural language input in order to derive a desired action. The SLU component **440** may output a most likely NLU hypothesis, or multiple NLU hypotheses associated with respective confidence or other scores (such as probability scores, etc.).

(67) As shown in FIG. **4**, the system **120** may include or otherwise communicate with one or more skills **425**. As noted above, a “skill” may refer to software, that may be placed on a machine or a virtual machine (e.g., software that may be launched in a virtual instance when called), configured to process NLU output data and perform one or more actions in response thereto. For example, for NLU output data including a <PlayMusic> intent, an “artist” entity type, and an artist name as an entity value, a music skill **425** may be called to output music sung by the indicated artist. For further example, for NLU output data including a <TurnOn> intent, a “device” entity type, and an entity value of “lights,” a smart home skill **425** may be called to cause one or more “smart” lights to operate in an “on” state. In another example, for NLU output data including an <OutputWeather> intent, a “location” entity type, and an entity value corresponding to a geographic location of the device **110**, a weather skill **425** may be called to output weather information for the geographic location. For further example, for NLU output data including a <BookRide> intent, a taxi skill **425** may be called to book a requested ride. In another example, for NLU output data including a <BuyPizza> intent, a restaurant skill **425** may be called to place an order for a pizza.

(68) A skill **425** may operate within the system **100**, e.g., as a component of the system **120**, the device **110**, a restaurant electronic ordering system, a taxi electronic booking system, etc., in order to complete certain functions. Inputs to a skill **425** may come from speech processing interactions or through other interactions or input sources.

(69) A skill **425** may be associated with a domain. A non-limiting list of example domains includes a smart home domain, a music domain, a video domain, a weather domain, a communications domain, a flash briefing domain, a shopping domain, and a custom domain.

(70) The system **120** may include a TTS component **480** that generates audio data including synthesized speech. The data input to the TTS component **480** may come from a skill **425**, the orchestrator component **430**, or another component of the system **120**.

(71) In one method of synthesis called “unit selection,” the TTS component **480** may match input data against a database of recorded speech. The TTS component **480** may select matching units of recorded speech and concatenate the units together to form audio data. In another method of synthesis called “parametric synthesis,” the TTS component **480** may vary parameters such as frequency, volume, and noise to determine audio data including an artificial speech waveform. Parametric synthesis may use a computerized voice generator, sometimes called a vocoder. In another technique, TTS may rely on neural networks or other machine learning components to process text data into audio data or the like to be processed by a vocoder and/or output as audio of synthesized speech.

(72) The system **120** may include a user recognition component **495**. The user recognition component **495** may recognize one or more users using various data. The user recognition component **495** may take as input the audio data **411**. The user recognition component **495** may perform user recognition by comparing speech characteristics in the audio data **411** to stored speech characteristics of users. The user recognition component **495** may additionally or alternatively perform user recognition by comparing biometric data (e.g., fingerprint data, iris data, retina data, etc.), received by the system **120** in correlation with a natural language input, to stored biometric data of users. The user recognition component **495** may additionally or alternatively perform user recognition by comparing image data (e.g., including a representation of one or more features of a user), received by the system **120** in correlation with a natural language input, with stored image data including representations of features of different users. The user recognition component **495** may also perform other or additional user recognition processes. For a

particular natural language input, the user recognition component **495** may perform processing with respect to stored data of users associated with the device **110** that received the natural language input. (73) The user recognition component **495** may determine whether a natural language input originated from a particular user. For example, the user recognition component **495** may determine a first value representing a likelihood that a natural language input originated from a first user, a second value representing a likelihood that the natural language input originated from a second user, etc. The user recognition component **495** may also determine an overall confidence regarding the accuracy of user recognition processing.

(74) The user recognition component **495** may output a single user identifier corresponding to the most likely user that originated the natural language input. Alternatively, the user recognition component **495** may output multiple user identifiers (e.g., in the form of an N-best list) with respective values representing likelihoods of respective users originating the natural language input. The output of the user recognition component **495** may be used to inform NLU processing, processing performed by a skill **425**, and/or processing performed by other components of the system **120** and/or other systems.

(75) The system **120** may include profile storage **470**. The profile storage **470** may include a variety of data related to individual users, groups of users, devices, etc., that interact with the system **120**. As used herein, a “profile” refers to a set of data associated with a user, group of users, device, etc. The data of a profile may include preferences specific to the user, group of users, device, etc.; input and output capabilities of one or more devices; internet connectivity data; user bibliographic data; subscription data; skill enablement data; and/or other data.

(76) The profile storage **470** may include one or more user profiles. Each user profile may be associated with a different user identifier. Each user profile may include various user identifying data (e.g., name, gender, address, one or more languages, etc.). Each user profile may also include preferences of the user. Each user profile may include one or more device identifiers, each representing a respective device registered to the user. Each user profile may include skill identifiers that identify the skills **425** that the user has enabled. When a user enables a skill **425**, the user is providing the system **120** with permission to allow the skill **425** to execute with respect to the user's natural language inputs. If a user does not enable a skill **425**, the system **120** may not execute the skill **425** with respect to the user's natural language inputs.

(77) The profile storage **470** may include one or more group profiles. Each group profile may be associated with a different group identifier. A group profile may be specific to a group of users. That is, a group profile may be associated with two or more individual user profiles. For example, a group profile may be a household profile that is associated with user profiles associated with multiple users of a single household. A group profile may include preferences shared by all the user profiles associated therewith. Each user profile associated with a group profile may additionally include preferences specific to the user associated therewith. That is, a user profile may include preferences unique from one or more other user profiles associated with the same group profile. A user profile may be a stand-alone profile or may be associated with a group profile. A group profile may be associated with (or include) one or more device profiles corresponding to one or more devices associated with the group profile.

(78) The profile storage **470** may include one or more device profiles. Each device profile may be associated with a different device identifier. A device profile may include various device identifying data, input/output characteristics, networking characteristics, etc. A device profile may also include one or more user identifiers, corresponding to one or more user profiles associated with the device profile. For example, a household device's profile may include the user identifiers of users of the household.

(79) The system **120** may also include a dialog management component **485**, which may manage various aspects of an ongoing dialog between the user **5** and the system **100**. Example components that may be included within the dialog management component **485**, as well as example operations that may be performed by such components, are described in detail below in connection with FIGS. **6-8**. As described in more detail below, in some implementations, the dialog management component **485** may include one or more components of the humor management component **130** described above.

(80) The system **120** may include a sentiment detection component **435** configured to analyze image data representing a face of the user **5**, and/or speech of the user (in particular tone, words, used, etc.), to determine a sentiment (e.g., happy, sad, mad, etc.) of the user **5**. Various processing described herein may

be based on the sentiment. For example, in some implementations, a determination to invoke the humor generator **150** to generate a pun based on one or more keywords indicative of a current context of a dialog between the user **5** and the system **100**, or otherwise, may be based, at least in part on an output of the sentiment detection component **435**. For example, if sentiment data from the sentiment detection component **435** indicates a user is frustrated or rushed, the orchestrator component **430**, dialog management component **485**, or other component may cause the humor generator **150** to be temporarily disabled and/or to ignore the output of the humor generator **150** for purposes of an exchange with a user. Further, if sentiment data from the sentiment detection component **435** indicates a user is happy or in a good mood, the orchestrator component **430**, dialog management component **485**, or other component may cause the humor generator **150** to be enabled and/or to be more likely to use the output of the humor generator **150** for purposes of an exchange with a user.

(81) The foregoing describes illustrative components and processing of the system **120**. With reference to FIG. 5, the following describes illustrative components and processing of the device **110**. As noted previously in connection with FIG. 4, in some embodiments, the system **120** may receive the audio data **411** from the device **110**, to recognize speech corresponding to a spoken natural language in the received audio data **411**, and to perform functions in response to the recognized speech. In some embodiments, these functions may involve sending directives (e.g., commands) from the system **120** to the device **110** to cause the device **110** to perform an action, such as to output synthesized speech (responsive to the spoken natural language input) via a loudspeaker(s), and/or to control one or more secondary devices by sending control commands to the one or more secondary devices. In other embodiments the device may perform various speech processing operations on its own and/or in conjunction with the system **120**.

(82) Thus, when the device **110** is able to communicate with the system **120** over the network(s) **199**, some or all of the functions capable of being performed by the system **120** may be performed by sending one or more directives over the network(s) **199** to the device **110**, which, in turn, may process the directive(s) and perform one or more corresponding actions. For example, the system **120**, using a remote directive that is included in response data (e.g., a remote response), may instruct the device **110** to output synthesized speech via a loudspeaker(s) of (or otherwise associated with) the device **110**, to output content (e.g., music) via the loudspeaker(s) of (or otherwise associated with) the device **110**, to display content on a display of (or otherwise associated with) the device **110**, and/or to send a directive to a secondary device (e.g., a directive to turn on a smart light). It will be appreciated that the system **120** may be configured to provide other functions in addition to those discussed herein, such as, without limitation, providing step-by-step directions for navigating from an origin location to a destination location, conducting an electronic commerce transaction on behalf of the user **5** as part of a shopping function, establishing a communication session (e.g., an audio or video call) between the user **5** and another user, and so on.

(83) As noted above, in some implementations, the device **110** may include a wakeword detection component **420** configured to detect a wakeword (e.g., “Alexa”) that indicates to the device **110** that the audio data **411** is to be processed for determining NLU output data. In some embodiments, a hybrid selector **524** of the device **110** (shown in FIG. 5) may send the audio data **411** to the wakeword detection component **420**. If the wakeword detection component **420** detects a wakeword in the audio data **411**, the wakeword detection component **420** may send an indication of such detection to the hybrid selector **524**. In response to receiving the indication, the hybrid selector **524** may send the audio data **411** to the system **120** and/or an on-device ASR component **550**. The wakeword detection component **420** may also send an indication, to the hybrid selector **524**, that a wakeword was not detected. In response to receiving such an indication, the hybrid selector **524** may refrain from sending the audio data **411** to the system **120**, and may prevent the on-device ASR component **550** from processing the audio data **411**. In this situation, the audio data **411** can be discarded.

(84) The device **110** may conduct its own speech processing using on-device language processing components (such as an on-device SLU component **540**, an on-device ASR component **550**, and/or an on-device NLU component **560**) similar to the manner discussed above with respect to the speech processing system-implemented SLU component **440**, ASR component **450**, and NLU component **460**. The device **110** may also internally include, or otherwise have access to, other components such as one or more skills **525** (configured to operate in a similar manner as the system-implemented skills **425**), a user recognition

component **595** (configured to operate in a similar manner as the system-implemented user recognition component **495**), profile storage **570** (configured to store similar profile data as the system-implemented profile storage **470**), a dialog management component **585** (configured to operate in a similar manner as the system-implemented dialog management component **485**), a sentiment detection component **535** (configured to operate in a similar manner as the system-implemented sentiment detection component **435**), a TTS **580** (configured to operate in a similar manner as the system-implemented TTS **480**), and other components. As described in more detail below, in some implementations, the dialog management component **585** of the device **110** may include one or more components of the humor management component **130** described above. In at least some embodiments, the on-device profile storage **570** may store profile data only for a user or group of users specifically associated with the device **110**.

(85) In some embodiments, the on-device language processing components may not have the same capabilities as the language processing components implemented by the system **120**. For example, the on-device language processing components may be configured to handle only a subset of the natural language inputs that may be handled by the speech processing system-implemented language processing components. For example, such subset of natural language inputs may correspond to local-type natural language inputs, such as those controlling devices or components associated with a user's home. In such circumstances, the on-device language processing components may be able to interpret and respond to a local-type natural language input more quickly than processing that involves the system **120**. If the device **110** attempts to process a natural language input for which the on-device language processing components are not necessarily best suited, the NLU output data, determined by the on-device components, may have a low confidence or other metric indicating that the processing by the on-device language processing components may not be as accurate as the processing that can be done by the system **120**.

(86) The hybrid selector **524** of the device **110** may include a hybrid proxy (HP) **526** configured to proxy traffic to/from the system **120**. For example, the HP **526** may be configured to send messages to/from a hybrid execution controller (HEC) **527** of the hybrid selector **524**. For example, command/directive data received from the system **120** can be sent to the HEC **527** using the HP **526**. The HP **526** may also be configured to allow the audio data **411** to pass to the system **120** while also receiving (e.g., intercepting) this audio data **411** and sending the audio data **411** to the HEC **527**.

(87) In some embodiments, the hybrid selector **524** may further include a local request orchestrator (LRO) **528** configured to notify the on-device ASR component **550** about the availability of the audio data **411**, and to otherwise initiate the operations of on-device language processing when the audio data **411** becomes available. In general, the hybrid selector **524** may control execution of on-device language processing, such as by sending “execute” and “terminate” events/instructions. An “execute” event may instruct a component to continue any suspended execution (e.g., by instructing the component to execute on a previously determined intent in order to determine a directive). Meanwhile, a “terminate” event may instruct a component to terminate further execution, such as when the device **110** receives directive data from the system **120** and chooses to use that remotely determined directive data.

(88) Thus, when the audio data **411** is received, the HP **526** may allow the audio data **411** to pass through to the system **120** and the HP **526** may also input the audio data **411** to the on-device ASR component **550** by routing the audio data **411** through the HEC **527** of the hybrid selector **524**, whereby the LRO **528** notifies the on-device ASR component **550** of the audio data **411**. At this point, the hybrid selector **524** may wait for response data from either or both the system **120** and/or the on-device language processing components. However, the disclosure is not limited thereto, and in some examples the hybrid selector **524** may send the audio data **411** only to the on-device ASR component **550** without departing from the disclosure. For example, the device **110** may process the audio data **411** on-device without sending the audio data **411** to the system **120**.

(89) The on-device ASR component **550** may be configured to receive the audio data **411** from the hybrid selector **524**, and to recognize speech in the audio data **411**, and the on-device NLU component **560** may be configured to determine an intent from the recognized speech (an optionally one or more named entities), and to determine how to act on the intent by generating NLU output data that may include directive data (e.g., instructing a component to perform an action). In some cases, a directive may include a description of the intent (e.g., an intent to turn off {device A}). In some cases, a directive may include

(e.g., encode) an identifier of a second device(s), such as kitchen lights, and an operation to be performed at the second device(s). Directive data may be formatted using Java, such as JavaScript syntax, or JavaScript-based syntax. This may include formatting the directive using JSON. In some embodiments, a device-determined directive may be serialized, much like how remotely-determined directives may be serialized for transmission in data packets over the network(s) **199**. In some embodiments, a device-determined directive may be formatted as a programmatic application programming interface (API) call with a same logical operation as a remotely-determined directive. In other words, a device-determined directive may mimic a remotely-determined directive by using a same, or a similar, format as the remotely-determined directive.

(90) A NLU hypothesis (output by the on-device NLU component **560**) may be selected as usable to respond to a natural language input, and local response data may be sent (e.g., local NLU output data, local knowledge base information, internet search results, and/or local directive data) to the hybrid selector **524**, such as a “ReadyToExecute” response. The hybrid selector **524** may then determine whether to use directive data from the on-device components to respond to the natural language input, to use directive data received from the system **120**, assuming a remote response is even received (e.g., when the device **110** is able to access the system **120** over the network(s) **199**), or to determine output data requesting additional information from the user **5**.

(91) The device **110** and/or the system **120** may associate a unique identifier with each natural language input. The device **110** may include the unique identifier when sending the audio data **411** to the system **120**, and the response data from the system **120** may include the unique identifier to identify to which natural language input the response data corresponds.

(92) In some embodiments, the device **110** may include one or more skills **525** that may operate similar to the system-implemented skill(s) **425** described above. The skill(s) **525** installed on (or in communication with) the device **110** may include, without limitation, a smart home skill and/or a device control skill configured to control a second device(s), a music skill configured to output music, a navigation skill configured to output directions, a shopping skill configured to conduct an electronic purchase, and/or the like.

(93) FIG. **6** shows an example implementation of the dialog management component **485/585**, which may be implemented by the system **120** (as shown in FIG. **4**), by the device **110** (as shown in FIG. **5**), or elsewhere in the system **100**. In some embodiments, the dialog management component **485/585** may be implemented as a skill **425/525**, or as a component of a skill **425/525**. As shown in FIG. **6**, in some implementations, in addition to the humor management component **130** (described above), the dialog management component **485/585** may include a dialog orchestrator component **610**, a response management component **620**, and a dialog history storage **630**.

(94) The dialog orchestrator component **610** may be configured to receive natural language input data **605**, e.g., data corresponding to a natural language input provided by the user **5**. The natural language input data **605** may include, for example, ASR output data (e.g., data representing words) and/or NLU output data (e.g., data representing an intent, entity, command, etc.) corresponding to received audio data **411** generated in response to an utterance by the user **5**. The dialog orchestrator component **610** may be configured to coordinate the transmission of data between components of the dialog management component **485/585**. In some implementations, for example, the dialog orchestrator component **610** may selectively call one or more components of the dialog management component **485/585** based on a determined intent of the natural language input data **605**. The response management component **620** may determine response data **615** that is to be sent to the device **110** or another component of the system **100**.

(95) For example, in some implementations, the response data **615** may cause the device **110** to output audio corresponding to the response data **615** (e.g., by processing the response data **615** with a TTS component **480/580**), or to output text corresponding to a system-generated response, e.g., via a display of the device **110**. In some implementations, the response data **615** may additionally or alternatively cause the device **110** to perform an operation (e.g., to begin playing music), or may cause another device to take a certain action (e.g., to cause a “smart light” to turn on or off). The response management component **620** may determine and/or generate such response data **615** in any of a number of ways. As but a few examples, the response management component **620** may (A) answer a question corresponding to the spoken natural language input data **605**, e.g., by retrieving responsive data from a knowledgebase,

(B) retrieve data from an information source (e.g., a weather application) to respond to a command (e.g., “tell me about today's weather”), (C) retrieve requested content (e.g., a song or story) from a datastore, (D) determine a command to provide to a home automation system (e.g., a command to turn on the living room lights), (E) initiate a skill **425/525** (e.g., to begin playing Jeopardy), (F) tell a joke about a particular topic in response to a command (e.g., “tell me a joke about chemistry,” (G) output a pun relating to a particular subject in response to a command (e.g., “tell me a pun relating to lawyers”), etc.

(96) At runtime, the system **120**/device **110** may receive natural language input data **605** corresponding to a dialog. As used herein, a “dialog” may refer to an exchange of related natural language input data **605** and system-generated response data **615**. A dialog may be goal-oriented, meaning the dialog is directed to the performance of a specific action (e.g., figuring out what music the system **100** should play). Receipt of natural language input data **605** and performance of a corresponding action (i.e., output of a system-generated response) may be referred to as a dialog “turn.” A dialog identifier may be associated with multiple related turns corresponding to consecutive related natural language inputs and system responses. Each turn may be associated with a respective turn identifier. One natural language input may be considered related to a subsequent natural language input, thereby causing a single dialog identifier to be associated with both natural language inputs. A first natural language input may be considered related to a second (subsequent) natural language input based on, for example, a length of time between receipt of the first and second natural language inputs, a length of time between performance of a system-generated response to the first natural language input and receipt of the second natural language input, the similarity of the subject matter of the first and second natural language inputs, and/or the similarity of the subject matter of the second natural language input and the system-generated response to the first natural language input.

(97) As disclosed above, the NLU component **460/560** may be configured to determine an intent of a natural language input. At runtime, the NLU component **460/560** may determine that a first natural input language corresponds to an intent associated with the dialog management component **485/585**, e.g., a skill **425/525**. In response to such a determination, first natural language input data **605** may be sent to the dialog management component **485/585**, resulting in the dialog management component **485/585** becoming “in focus” for a dialog including the first natural language input data **605**. The dialog orchestrator component **610** of the dialog management component **485/585** may send the first natural language input data **605** to the response management component **620** for processing to determine response data **615** to send to the device **110** or elsewhere. Thereafter, when a second natural language input is received, the system **120**/device **110** may determine that the second natural language input corresponds to the same dialog as the first natural language input, and thus determine that the dialog management component **485/585** remains in focus for the dialog. Based on the dialog management component **485/585** being in focus, second natural language input data **605** corresponding to the second natural language input may not undergo NLU processing by the NLU component **460/560**. Rather, the second natural language input data **605**, e.g., including ASR output data corresponding to the second natural language input, may be sent to the dialog management component **485/585**.

(98) In the foregoing situation, an intent classifier of the response management component **620** may process the ASR output data to determine an intent of the second natural language input data **605**. The intent classifier of the response management component **620** may be configured to determine one or more intents executable by the dialog management component **485/585**. In some embodiments, the dialog management component **485/585** may, for example, be configured to engage the user **5** in a storytelling user experience that optionally includes question and answering functionality. In such embodiments, the intent classifier may be configured to determine intents such as a <RequestStory> intent indicating that natural language input data **605** request that a story be output; a <StoryNavigation> intent indicating that natural language input data **605** requests a different section (e.g., previous page, next page, etc.) of a book be output; a <QuestionAnswer> intent indicating that natural language input data **605** requests an answer to a question represented in the natural language input data **605**; a <ChangeStory> intent indicating that natural language input data **605** requests a different story be output; and/or a <Conversation> intent indicating that natural language input data **605** provides information requested by the dialog management system **485/585** in order to perform an action.

(99) The intent classifier of the response management component **620** may be configured using various

techniques. In some embodiments, the intent classifier may implement a machine learning (ML) model(s). Various machine learning techniques may be used to train and operate ML models. A ML model may be trained and operated according to various ML techniques. Such techniques may include, for example, neural networks (such as deep neural networks and/or recurrent neural networks), inference engines, trained classifiers, etc. Examples of trained classifiers include Support Vector Machines (SVMs), neural networks, decision trees, AdaBoost (short for “Adaptive Boosting”) combined with decision trees, and random forests. Focusing on SVM as an example, SVM is a supervised learning model with associated learning algorithms that analyze data and recognize patterns in the data, and which are commonly used for classification and regression analysis. Given a set of training examples, each marked as belonging to one of two categories, an SVM training algorithm builds a model that assigns new examples into one category or the other, making it a non-probabilistic binary linear classifier. More complex SVM models may be built with the training set identifying more than two categories, with the SVM determining which category is most similar to input data. An SVM model may be mapped so that the examples of the separate categories are divided by clear gaps. New examples are then mapped into that same space and predicted to belong to a category based on which side of the gaps they fall on. Classifiers may issue a “score” indicating which category the data most closely matches. The score may provide an indication of how closely the data matches the category.

(100) In order to apply the machine learning techniques, the machine learning processes themselves need to be trained. Training a machine learning component such as, in this case, one of the first or second models, requires establishing a “ground truth” for the training examples. In machine learning, the term “ground truth” refers to the accuracy of a training set's classification for supervised learning techniques. Various techniques may be used to train the models including backpropagation, statistical learning, supervised learning, semi-supervised learning, stochastic learning, or other known techniques.

(101) At runtime, inputs to the intent classifier of the response management component **620** may include, for example, a dialog identifier of the present dialog, a user identifier (e.g., output by the user recognition component **495/595**), the natural language input data **605**, user recognition data, sentiment data, and/or dialog state data (e.g., representing the presently executing functionality of the dialog management component **485/585**, such as storytelling or question and answering).

(102) In some embodiments, the dialog management component **485/585** may not include an intent classifier. In such embodiments, even when the dialog management component **485/585** is in focus, the NLU component **460/560** may be used to determine an intent of a subsequently received natural language input, and the natural language input data **605** received by the dialog management component **485/585** in such circumstance may include NLU output data in addition to or instead of ASR output data.

(103) In some embodiments, the dialog management component **485/585** may include a NER component (not illustrated). The NER component may process similarly to the NER component of the NLU component **460/560** described above. Such an NER component of the dialog management component **485/585** may be configured to determine entity types and values corresponding to the dialog management component **485/585**, and not necessarily all entity types and values determinable by the NER component of the NLU component **460/560**. In some embodiments, the NER component, of the dialog management component **485/585**, may implement a ML model.

(104) The dialog history storage **630** of the dialog management component **485/585** may store various data relating to one or more dialogs. For example, for a given dialog, the dialog history storage **630** may associate a dialog identifier with the natural language input data **605** of the dialog, the intent(s) determined for the natural language input(s) of the dialog, the natural language input data **605** received by the dialog management component **485/585**, and the response data **615** sent to the device **110** or elsewhere as part of the dialog. The dialog history storage **630** may receive data from the dialog orchestrator component **610** and/or the response management component **620**. Similarly, the intent classifier of the response management component **620** may use dialog history data, stored in the dialog history storage **630** for the one or more dialogs, to improve intent classification processing (i.e., to make it more likely that the intent determined by the intent classifier accurately represents the actual intent of the natural language input). Likewise, the humor management component **130** may use dialog history data, stored in the dialog history storage **630** for one or more dialogs, to determine whether a given dialog includes humorous content (e.g., by recognizing the presence of one or more puns in one or more

sentences of the dialog) and/or to represent one or more keywords representing the current context of one or more dialogs for use in generating a pun that make sense in that context.

(105) With respect to recognizing and/or inserting humor in a dialog, in some implementations, the humor recognizer **140** of the humor management component **130** may continuously monitor the text of one or more sentences of the dialog, such as by continuously feeding a set of tokens representing such text to the input of the pun explanation model **202** shown in FIG. 2A, together with the prefix that was used during the T5 model fine-tuning process, e.g., as described above in connection with FIG. 2C. The output of the pun classification model **204** of the humor recognizer **140** may then be provided to the response management component **620** for use in determining appropriate response data **615** when the output of the pun classification model **204** indicates that a pun has been recognized in the text of the dialog. For instance, as alluded to above, if the user **5** says something that either alone or together with previously output response data **615** is recognized as a pun by the humor recognizer **140**, the response management component **620** may cause the device **110** to output “I see what you did there!” or some similar phrase, to output a giggling or laughing sound, etc. Further, as discussed in more detail below, in some implementations, the response management component **620** may evaluate various inputs to determine whether to request that the humor generator **150** generate a pun (e.g., to send to the device **110** as response data **615**) that makes sense in the current context of the dialog. In some implementations, one of the inputs the response management component **620** may evaluate for that purpose may be the output of the humor recognizer **140**. Such an evaluation may enable the response management component **620** to inject humor into a dialog at an appropriate time, as the recognition of a pun in the current dialog may indicate that the user **5** is in a playful or joking mood.

(106) As noted above, the response management component **620** may determine appropriate responses to natural language inputs from the user **5**. In some implementations, the response management component **620** may determine, based on the current contextual state of the system **100**, that it would be a good time to cause the device **110** to output a humorous statement, such as a pun. In such a scenario, the response management component **620** may send a request to the humor management component **130** to generate a pun that can be output as response data **615**. As noted above, to make such a determination, the response management component **620** may receive inputs from various sources, including the output of humor recognizer **140** of the humor management component **130**. In some implementations, the response management component **620** may additionally or alternatively interact with the sentiment detection component **435/535** to determine the current sentiment of the user **5**, such as whether the user is happy, smiling, laughing, etc., or perhaps whether the user **5** is sad, frowning, or crying, and could perhaps use some cheering up. The determination by the response management component **620** of whether to request the generation of a pun by the humor generator **150** may further be based at least in part on the determined sentiment of the user **5**. Further, in some implementations, the response management component **620** may additionally or alternatively query the profile storage **470/570** for profile data corresponding to the user **5**, such as data indicating whether the user likes or dislikes receiving humorous responses from the system **100**. The determination by the response management component **620** of whether to request the generation of a pun by the humor generator **150** may further be based at least in part on such profile data.

(107) In some implementations, the profile data stored in the profile storage **470/570** may additionally include information such as the age of the user **5**, a geographic location of the user **5** or the device **110**, one or more languages spoken by the user **5**, one or more joke type preferences of the user **5**, one or more subject matter preference (e.g., cats, dogs, fishing, etc.) of the user **5**. In such implementations, the response management component **620** may send such additional information to the humor generator **150**, or the humor generator **150** may separately access such additional information, and the humor generator **150** may use it to generate humorous text, e.g., a pun, that is more closely related to the user's preferences and/or demographic data. In some implementations, for instance, one or more items of such preference data may be used to determine one or more of the keywords that are fed to the pun word classification model **304** shown in FIG. 3A.

(108) As discussed above in connection with FIG. 3A, the humor generator **150** of the humor management component **130** may include a keyword extractor **302** that is configured to identify one or more keywords that are to be fed to the pun word classification model **304**. In some implementations, in

response to receiving a request to generate a pun, the humor generator **150** may evaluate the content of the dialog history storage **630** to identify one or more sentences corresponding to one or more recent turns of one or more dialogs with the user **5**. The humor generator **150** may then use a keyword extraction algorithm, such as the RAKE algorithm, to extract keywords from such sentences. In some implementations, one or more keywords corresponding to one or more user preferences, may also be identified and used to supplement the keyword(s) extracted from the dialog turns. As an example, if the profile of the user **5** indicates that the user has an affinity for cats, the word “cat” can be at least occasionally be inserted into the list of keyword(s) extracted from content of the dialog history storage **630**.

(109) No matter how the keyword(s) are determined, as described above in connection with FIG. **3A**, the identified keyword(s) maybe fed to the pun classification model **304**, as well as the pun generation model **308**, to initiate the generation of a pun by the pun generation model **308**, provided the pun classification model **304** is able to identify a pun word pair that is sufficiently contextually related to the identified keyword(s), e.g., by having a confidence value that exceeds a threshold.

(110) Although not illustrated in FIG. **3A**, it should be appreciated that, in some implementations, the pun word selector **306** may further have access to a storage medium that stores data representing the senses of the various pun words and alternate words the pun word classification model **304** is trained to classify. In such implementations, the pun word selector **306** may retrieve such sense information, e.g., from such storage medium, and provide that sense information to the pun generation model **308**, together with an identified pun word pair and the determined keyword(s), to enable the pun generation model **308** to use that sense information to generate an even more suitable, context-situated pun. In other implementations, pun word/alternate word sense information may not be used when generating puns, and the pun word selector **306** may send only a selected pun word pair to the pun generation model **308** for use in generating a pun that make sense in the context of determined keyword(s).

(111) FIG. **7** shows an example process **700** that may be performed by one or more components of the dialog management component **485/585** to recognize and respond to humor during a dialog. In some implementations, for example, the device **110** (described above) may be a voice-controlled device, such as an Amazon Echo, and the user **5** may speak one or utterances in a vicinity of the device **110** during a dialog that is being managed by the dialog management component **485/585**. In other implementation, the device **110** may determine text corresponding to a dialog in other ways, such as in response to the user **5** typing on a keyboard of the device **110**.

(112) As shown in FIG. **7**, at a step **702** of the process **700**, the dialog management component **485/585** may receive a natural language input from a user device **110**. For example, as shown in FIG. **6**, natural language input data **605** representing an utterance (e.g., ASR data) by the user **5** may be received by the dialog orchestrator component **610**.

(113) At a step **704** of the process **700**, the dialog management component **485/585** may determine text corresponding to a dialog including the natural language input received at the step **702**. In some implementations, for example, the humor recognizer **140** of the humor management component **130** may determine the text of one or more sentences of a dialog represented by data stored in dialog history storage **630**. Such text may include a textual representation of the natural language input received at the step **702**, and may also include a textual representation of one or more responses (e.g., corresponding to response data **615**) sent to the device **110** before and/or after receipt of the natural language input.

(114) At a step **706** of the process **700**, the humor recognizer **140** may determine a pun explanation prefix that can be provided to the pun explanation model **202** (see FIG. **2A**), together with the text determined at the step **704**, to prompt the pun explanation model **202** to generate an explanation concerning the text. The pun explanation prefix determined at the step **706** may be the same prefix that was used to fine-tune a BERT model to generate the pun explanation model **202**, as described above in connection with FIG. **2C**.

(115) At a step **708** of the process **700**, the humor recognizer **140** may provide the text determined at the step **704** and the prefix determined at the step **706** to the pun explanation model **202**.

(116) At a step **710** of the process **700**, the pun explanation model **202** may process the text and prefix to generate an explanation for the text.

(117) At a step **712** of the process **700**, the humor recognizer **140** may provide the text (determined at the

step **704**) and the explanation (generated at the step **710**) to the pun classification model **204** (see FIG. 2A).

(118) At a step **714** of the process **700**, the pun classification model **204** may process the text and the generated explanation to determine whether or not the text represents a pun. As shown in FIG. 2A, for example, in some implementations, the pun classification model **204** may output a “pun indicator.” In some implementations, the pun indicator may be a Boolean value indicating whether or not the text represents a pun. In other implementation, the pun indicator may additionally or alternatively represent confidence value indicating a level of confidence that the text represents a pun.

(119) At a step **716** of the process **700**, the humor recognizer **140** may use the output of the pun classification model **204** to determine that the text represents a pun, such as by determining that a Boolean output has a certain value, e.g., “1,” or by determining that an output confidence value exceeds a threshold.

(120) At a step **718** of the process **700**, in response to determining that the text represents a pun, the humor recognizer **140** may send an indication of an identified pun to the response management component **620** (see FIG. 6).

(121) At a step **720** of the process **700**, the response management component **620** may generate response data **615** corresponding to the pun. For example, as noted above, in response to the humor recognizer **140** notifying the response management component **620** that something the user said, either alone or together with an output provided by the device **110**, represents a pun, the **620** may generate response data **615** to cause the device **110** to output “I see what you did there!,” to output a laughing or giggling sound, etc.

(122) At a step **722** of the process **700**, the response management component **620** may send the generated response data **615** to the device **110**, thus causing the device **110** to output a corresponding response. In some implementations, the generated response data **615** may comprise audio data that causes the device **110** to output an audio response. In other implementations, the generated response data **615** may be text data and a TTS component **480/580** may be used to convert that text data into audio data that can be sent to the device **110**. In still other implementations, the response management component **620** may additionally or alternatively send generated text data to the device **110** for display on a display screen of the device **110**.

(123) FIG. 8 shows an example process **800** that may be performed by one or more components of the dialog management component **485/585** to inject context-based humor into a dialog at an appropriate time. As noted above, in some implementations, the response management component **620** (shown in FIG. 6) may be responsible for determining, based on one or more inputs, whether to send a request to the humor generator **150** to generate a pun that makes sense in the current context of the dialog. Examples of inputs that maybe evaluated by the response management component **620** for that purpose are reflected in the steps **802**, **804**, **806** and **808** of the process **800**. As shown in FIG. 8, per the step **802**, the response management component **620** may receive a first indication of an identified pun (e.g., an output of the pun classification model **204** shown in FIG. 2A) from the humor recognizer **140**. Per the step **804**, the response management component **620** may additionally or alternatively receive a second indication of a detected sentiment of a user **5** operating the device **110**, such as whether the person is happy, smiling, or laughing. Per the step **806**, the response management component **620** may additionally or alternatively receive dialog history data from the dialog history storage **630**, such as the text of one or more recent dialogs with the user **5** of the device **110**. Per the step **808**, the response management component **620** may additionally or alternatively receive profile data corresponding to the user **5** of the device **110**.

(124) At a step **810** of the process **800**, the response management component **620** may determine, based on one or more of the foregoing inputs (or otherwise), to cause the device **110** to output humorous content. For example, if one or more of such inputs indicate that the user **5** is likely in a jovial mood, the response management component **620** may determine that it is an appropriate time to cause the device **110** to output audio that is intended to make the user **5** laugh.

(125) In response to determining to cause the device **110** to output humorous content (per the step **810**), the response management component **620** may (per a step **812**) send a request to the humor generator **150** to generate a pun that makes sense given the current context of the dialog with the user **5**.

(126) At a step **814** of the process **800**, the keyword extractor **302** of the humor generator **150** (see FIG. 3A) may determine one or more keywords corresponding to one or more dialogs involving the user **5**. For

example, as noted above, in some implementations, the keyword extractor **302** may employ the rapid automatic keyword extraction (RAKE) algorithm to extract one or more keywords from such one or more recent utterances spoken by the user **5** and/or one or more recent responses provided by the device **110** to the user **5**.

(127) At a step **816** of the process **800** the keyword extractor **302** may provide the extracted keyword(s) to the pun word classification model **304**.

(128) At a step **818** of the process **800**, the pun word classification model **304** may process the input keyword(s) to determine confidence values indicating the suitability of respective pun words pair to generate a pun corresponding to such keyword(s). If none of the confidence values exceeds a threshold value, the humor generator **150** may send a message to the response management component **620** indicating that a suitable pun cannot be generated given the current contextual circumstances.

(129) At a step **820** of the process **800**, if at least one confidence value determined by the pun word classification model **304** exceeds the threshold value, the pun word selector **306** may evaluate the confidence values output by the pun word classification model **304** to identify the pun word pair having the highest confidence value.

(130) At a step **822** of the process **800**, the pun word selector **306** may (optionally) determine senses (i.e., meanings) for the pun word and the alternate word in the pun word pair identified per the step **820**.

(131) At a step **824** of the process **800**, the humor generator **150** may determine a prefix to use to instruct the pun generation model **308** to generate a pun based on its various inputs (e.g., the same prefix that was used to fine-tune the pun generation model **308** as described in connection with FIG. 3D).

(132) At a step **826** of the process **800**, the humor generator **150** may provide the pun word pair identified at the step **820** and (optionally) the senses identified at the step **822** to the pun generation model **308**, together with the keyword(s) determined at the step **814** and the prefix determined at the step **824** to the pun generation model **308**.

(133) At a step **828** of the process **800**, the pun generation model **308** may, based on the foregoing inputs, generate a pun that makes sense in the context of the keyword(s) determined at the step **814**.

(134) At a step **830** of the process **800**, the pun word selector **306** of the humor generator **150** may send the generated pun to the response management component **620**. Alternatively, as noted previously, in some circumstances the humor generator **150** may instead send to the response management component **620** an indication that the humor generator **150** was unable to generate a pun, e.g., when none of the confidence values determined by the pun word classification model **304** exceeded the threshold value.

(135) At a step **832** of the process **800**, in the event that the response management component **620** receives a generated pun from the humor generator **150**, the response management component **620** may cause the device **110** to output a response representing the generated pun, such as by sending response data **615** representing the generated pun to the device **110**, or perhaps by sending response data **615** representing the generated pun a TTS components **480/580** to convert the response data **615** into audio data to send to the device **110** for output as an audio response. If, on the other hand, the response management component **620** received from the humor generator **150** and indication that a suitable pun could not be generated, the response management component **620** may instead determine not to cause the device **110** to output humorous content given the current contextual circumstances.

(136) FIG. 9 is a block diagram conceptually illustrating a device **110** that may be used with the system **120**. FIG. 10 is a block diagram conceptually illustrating example components of a remote device, such as the system **120** or a skill **425/525**. A system (**120/425/525**) may include one or more servers. A “server” as used herein may refer to a traditional server as understood in a server/client computing structure but may also refer to a number of different computing components that may assist with the operations discussed herein. For example, a server may include one or more physical computing components (such as a rack server) that are connected to other devices/components either physically and/or over a network and is capable of performing computing operations. A server may also include one or more virtual machines that emulates a computer system and is run on one or across multiple devices. A server may also include other combinations of hardware, software, firmware, or the like to perform operations discussed herein. The system (**120/425/525**) may be configured to operate using one or more of a client-server model, a computer bureau model, grid computing techniques, fog computing techniques, mainframe techniques, utility computing techniques, a peer-to-peer model, sandbox

techniques, or other computing techniques.

(137) Multiple systems (**120/425/525**) may be included in the system **100** of the present disclosure, such as one or more systems **120** and/or one or more skills **425/525**. In operation, each of these systems may include computer-readable and computer-executable instructions that reside on the respective device (**120/425/525**), as will be discussed further below.

(138) Each of these devices (**110/120/425/525**) may include one or more controllers/processors (**904/1004**), which may each include a central processing unit (CPU) for processing data and computer-readable instructions, and a memory (**906/1006**) for storing data and instructions of the respective device. The memories (**906/1006**) may individually include volatile random access memory (RAM), non-volatile read only memory (ROM), non-volatile magnetoresistive memory (MRAM), and/or other types of memory. Each device (**110/120/425/525**) may also include a data storage component (**908/1008**) for storing data and controller/processor-executable instructions. Each data storage component (**908/1008**) may individually include one or more non-volatile storage types such as magnetic storage, optical storage, solid-state storage, etc. Each device (**110/120/425/525**) may also be connected to removable or external non-volatile memory and/or storage (such as a removable memory card, memory key drive, networked storage, etc.) through respective input/output device interfaces (**902/1002**).

(139) Computer instructions for operating each device (**110/120/425/525**) and its various components may be executed by the respective device's controller(s)/processor(s) (**904/1004**), using the memory (**906/1006**) as temporary “working” storage at runtime. A device's computer instructions may be stored in a non-transitory manner in non-volatile memory (**906/1006**), storage (**908/1008**), or an external device(s). Alternatively, some or all of the executable instructions may be embedded in hardware or firmware on the respective device in addition to or instead of software.

(140) Each device (**110/120/425/525**) includes input/output device interfaces (**902/1002**). A variety of components may be connected through the input/output device interfaces (**902/1002**), as will be discussed further below. Additionally, each device (**110/120/425**) may include an address/data bus (**1624/1024**) for conveying data among components of the respective device. Each component within a device (**110/120/425/525**) may also be directly connected to other components in addition to (or instead of) being connected to other components across the bus (**1624/1024**).

(141) Referring to FIG. **9**, the device **110** may include input/output device interfaces **902** that connect to a variety of components such as an audio output component such as a speaker **912**, a wired headset or a wireless headset (not illustrated), or other component capable of outputting audio. The device **110** may also include an audio capture component. The audio capture component may be, for example, a microphone **920** or array of microphones, a wired headset or a wireless headset (not illustrated), etc. If an array of microphones is included, approximate distance to a sound's point of origin may be determined by acoustic localization based on time and amplitude differences between sounds captured by different microphones of the array. The device **110** may additionally include a display **916** for displaying content. The device **110** may further include a camera **918**.

(142) Via antenna(s) **914**, the input/output device interfaces **902** may connect to one or more networks **199** via a wireless local area network (WLAN) (such as Wi-Fi) radio, Bluetooth, and/or wireless network radio, such as a radio capable of communication with a wireless communication network such as a Long Term Evolution (LTE) network, WiMAX network, 3G network, 4G network, 5G network, etc. A wired connection such as Ethernet may also be supported. Through the network(s) **199**, the system may be distributed across a networked environment. The I/O device interface (**902/1002**) may also include communication components that allow data to be exchanged between devices such as different physical servers in a collection of servers or other components.

(143) The components of the device **110**, the system **120**, and/or the skill **425/525** may include their own dedicated processors, memory, and/or storage. Alternatively, one or more of the components of the device **110**, the system **120**, and/or the skill **425/525** may utilize the I/O interfaces (**902/1002**), processor(s) (**904/1004**), memory (**906/1006**), and/or storage (**908/1008**) of the device **110**, the system **120**, or the skill **425/525**, respectively. Thus, the ASR component **450/560** may have its own I/O interface(s), processor(s), memory, and/or storage; the NLU component **460/560** may have its own I/O interface(s), processor(s), memory, and/or storage; and so forth for the various components discussed herein.

(144) As noted above, multiple devices may be employed in a single system. In such a multi-device

system, each of the devices may include different components for performing different aspects of the system's processing. The multiple devices may include overlapping components. The components of the device **110**, the system **120**, and the skill **425/525**, as described herein, are illustrative, and may be located as a stand-alone device or may be included, in whole or in part, as a component of a larger device or system.

(145) As illustrated in FIG. **11**, multiple devices (**110a-110j**, **120**, **425/525**) may contain components of the system and the devices may be connected over a network(s) **199**. The network(s) **199** may include a local or private network or may include a wide network such as the Internet. Devices may be connected to the network(s) **199** through either wired or wireless connections. For example, a speech-controllable device **110a**, a smart phone **110b**, a smart watch **110c**, a tablet computer **110d**, a vehicle **110e**, a speech-controllable display device **110f**, a smart television **110g**, a washer/dryer **110h**, a refrigerator **110i**, and/or a microwave **110j** may be connected to the network(s) **199** through a wireless service provider, over a Wi-Fi or cellular network connection, or the like. Other devices are included as network-connected support devices, such as the system **120**, the skill **425/525**, and/or others. The support devices may connect to the network(s) **199** through a wired connection or wireless connection.

(146) The concepts disclosed herein may be applied within a number of different devices and computer systems, including, for example, general-purpose computing systems, speech processing systems, and distributed computing environments.

(147) The above aspects of the present disclosure are meant to be illustrative. They were chosen to explain the principles and application of the disclosure and are not intended to be exhaustive or to limit the disclosure. Many modifications and variations of the disclosed aspects may be apparent to those of skill in the art. Persons having ordinary skill in the field of computers and speech processing should recognize that components and process steps described herein may be interchangeable with other components or steps, or combinations of components or steps, and still achieve the benefits and advantages of the present disclosure. Moreover, it should be apparent to one skilled in the art, that the disclosure may be practiced without some or all of the specific details and steps disclosed herein.

(148) Aspects of the disclosed system may be implemented as a computer method or as an article of manufacture such as a memory device or non-transitory computer readable storage medium. The computer readable storage medium may be readable by a computer and may comprise instructions for causing a computer or other device to perform processes described in the present disclosure. The computer readable storage medium may be implemented by a volatile computer memory, non-volatile computer memory, hard drive, solid-state memory, flash drive, removable disk, and/or other media. In addition, components of system may be implemented as in firmware or hardware, such as an acoustic front end (AFE), which comprises, among other things, analog and/or digital filters (e.g., filters configured as firmware to a digital signal processor (DSP)).

(149) Conditional language used herein, such as, among others, “can,” “could,” “might,” “may,” “e.g.,” and the like, unless specifically stated otherwise, or otherwise understood within the context as used, is generally intended to convey that certain embodiments include, while other embodiments do not include, certain features, elements and/or steps. Thus, such conditional language is not generally intended to imply that features, elements, and/or steps are in any way required for one or more embodiments or that one or more embodiments necessarily include logic for deciding, with or without other input or prompting, whether these features, elements, and/or steps are included or are to be performed in any particular embodiment. The terms “comprising,” “including,” “having,” and the like are synonymous and are used inclusively, in an open-ended fashion, and do not exclude additional elements, features, acts, operations, and so forth. Also, the term “or” is used in its inclusive sense (and not in its exclusive sense) so that when used, for example, to connect a list of elements, the term “or” means one, some, or all of the elements in the list.

(150) Disjunctive language such as the phrase “at least one of X, Y, Z,” unless specifically stated otherwise, is understood with the context as used in general to present that an item, term, etc., may be either X, Y, or Z, or any combination thereof (e.g., X, Y, and/or Z). Thus, such disjunctive language is not generally intended to, and should not, imply that certain embodiments require at least one of X, at least one of Y, or at least one of Z to each be present.

(151) As used in this disclosure, the term “a” or “one” may include one or more items unless specifically

stated otherwise. Further, the phrase “based on” is intended to mean “based at least in part on” unless specifically stated otherwise.

Claims

1. A computer-implemented method comprising: receiving, from a first device, first audio data representing an utterance corresponding to a dialog between a user of the first device and a system; performing speech processing using the first audio data to determine first data representing a plurality of words in the dialog; processing the first data using a bidirectional encoder representations from transformers (BERT) model configured to use input data to classify a candidate word pair as suitable for generating a pun that is contextually related to the input data; using an output of the BERT model to select a first word pair suitable for generating a pun; processing the first word pair and the first data using a text-to-text transfer transformer (T5) model to generate second data representing a first pun that is contextually related to the first data; sending the second data to a dialog management component; sending, from the dialog management component to a speech synthesis component, the second data; processing the second data by the speech synthesis component to generate second audio data representing the first pun; and sending the second audio data to the first device.
2. The computer-implemented method of claim 1, wherein the T5 model generates the first pun based at least in part on a first meaning of a first word of the first word pair and a second meaning of a second word of the first word pair.
3. The computer-implemented method of claim 1, further comprising, prior to receiving the first audio data: training the BERT model using training data including third data representing one or more keywords, fourth data representing a second word pair, and an indication whether the second word pair is suitable for generating a pun contextually related to the one or more keywords, with the indication being used to supervise the training.
4. The computer-implemented method of claim 1, further comprising: training the T5 model using training data including at least third data representing a text prefix corresponding to pun generation, fourth data representing a second word pair, fifth data representing one or more keywords, and sixth data representing a second pun, with the sixth data being used to supervise the training.
5. A computer-implemented method, comprising: determining first data representing one or more first keywords; processing the first data by at least one first component to determine a first word pair suitable for generating a pun; processing the first word pair and the one or more first keywords using a first machine learning model to generate a first pun that is contextually related to the one or more first keywords, wherein the first machine learning model is trained using training data including at least: second data representing a text prefix corresponding to pun generation, third data representing a second word pair, fourth data representing one or more second keywords, and fifth data representing a second pun; and causing a first device to output a representation of the first pun.
6. The computer-implemented method of claim 5, wherein the first machine learning model generates the first pun based at least in part on a first meaning of a first word of the first word pair and a second meaning of a second word of the first word pair.
7. The computer-implemented method of claim 5, wherein the first machine learning model comprises a text-to-text transfer transformer model.
8. The computer-implemented method of claim 7, further comprising: training the text-to-text transfer transformer model using the training data wherein the fifth data is second pun, wherein the fifth data is used to supervise the training.
9. The computer-implemented method of claim 5, wherein the at least one first component comprises a second machine learning model that uses the first data to classify respective candidate word pairs as suitable for generating a pun that is contextually related to the first data.
10. The computer-implemented method of claim 9, wherein the at least one first component further comprises a pun word selector configured to select the first word pair based on an output of the second machine learning model.
11. The computer-implemented method of claim 5, further comprising: receiving, by a dialog management component, a first communication; determining the first data based on the first

communication; generating, by the dialog management component, a response representing the first pun; and sending, to the first device, a second communication corresponding to the response to cause the first device to output the representation of the first pun.

12. The computer-implemented method of claim 5, wherein: the first machine learning model is configured to generate an output pun based at least in part on one or more input keywords and an input word pair such that the output pun is contextually related to the one or more input keywords.

13. A system, comprising: at least one processor; and at least one memory comprising instructions that, when executed by the at least one processor, cause the system to: determine first data representing one or more first keywords; process the first data by at least one first component to determine a first word pair suitable for generating a pun; process the first word pair and the one or more first keywords using a first machine learning model to generate a first pun that is contextually related to the one or more first keywords, wherein the first machine learning model is trained using training data including at least: second data representing a text prefix corresponding to pun generation, third data representing a second word pair, fourth data representing one or more second keywords, and fifth data representing a second pun; receive, from the first machine learning model, sixth data representing the first pun; and cause, based on the sixth data, a first device to output a representation of the first pun.

14. The system of claim 13, wherein the at least one memory further comprises additional instructions that, when executed by the at least one processor, cause the system to: process a first meaning of a first word of the first word pair and a second meaning of a second word of the first word pair with the first machine learning model to generate the first pun.

15. The system of claim 13, wherein the first machine learning model comprises a text-to-text transfer transformer model.

16. The system of claim 15, wherein the at least one memory further comprises additional instructions that, when executed by the at least one processor, cause the system to: train the text-to-text transfer transformer model using the training data, wherein the fifth data is used to supervise the training.

17. The system of claim 13, wherein the at least one first component comprises a second machine learning model, and the at least one memory further comprises additional instructions that, when executed by the at least one processor, cause the system to: process the first data with the second machine learning model to classify respective candidate word pairs as suitable for generating a pun that is contextually related to the first data.

18. The system of claim 17, wherein the at least one first component further comprises a pun word selector, and the at least one memory further comprises additional instructions that, when executed by the at least one processor, cause the system to: use the pun word selector to select the first word pair based on an output of the second machine learning model.

19. The system of claim 13, wherein the at least one memory further comprises additional instructions that, when executed by the at least one processor, cause the system to: receive, by a dialog management component, a first communication; determine the first data based on the first communication; generate, by the dialog management component and based on the sixth data, a response representing the first pun; and send, to the first device, a second communication corresponding to the response to cause the first device to output the representation of the first pun.

20. The system of claim 19, wherein: the first machine learning model is configured to generate an output pun based at least in part on one or more input keywords and an input word pair such that the output pun is contextually related to the one or more input keywords.
