

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12400663
Kind Code	B2
Date of Patent	August 26, 2025
Inventor(s)	Pasko; Stanislaw Ignacy

Speech interface device with caching component

Abstract

A speech interface device is configured to receive response data from a remote speech processing system for responding to user speech. This response data may be enhanced with information such as a remote ASR result(s) and a remote NLU result(s). The response data from the remote speech processing system may include one or more cacheable status indicators associated with the NLU result(s) and/or remote directive data, which indicate whether the remote NLU result(s) and/or the remote directive data are individually cacheable. A caching component of the speech interface device allows for caching at least some of this cacheable remote speech processing information, and using the cached information locally on the speech interface device when responding to user speech in the future. This allows for responding to user speech, even when the speech interface device is unable to communicate with a remote speech processing system over a wide area network.

Inventors:	Pasko; Stanislaw Ignacy (Gdansk, PL)
Applicant:	Amazon Technologies, Inc. (Seattle, WA)
Family ID:	1000008781699
Assignee:	Amazon Technologies, Inc. (Seattle, WA)
Appl. No.:	18/425465
Filed:	January 29, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240249725 A1	Jul. 25, 2024

Related U.S. Application Data

continuation parent-doc US 17902519 20220902 US 11887604 child-doc US 18425465
continuation parent-doc US 17018279 20200911 US 11437041 20220906 child-doc US 17902519

Publication Classification

Int. Cl.: **G10L15/30** (20130101); **G10L15/18** (20130101); **H04L67/5683** (20220101)

U.S. Cl.:

CPC **G10L15/30** (20130101); **G10L15/18** (20130101); **H04L67/5683** (20220501);

Field of Classification Search

CPC: G10L (15/18); G10L (15/32); G10L (15/22); G10L (15/30); G10L (25/51); G10L (15/1815); H04L (67/2857)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6204763	12/2000	Sone	N/A	N/A
6226749	12/2000	Carloganu	726/2	G06F 21/31
6408272	12/2001	White et al.	N/A	N/A
6836760	12/2003	Bellegarda	704/238	G10L 15/1822
8364694	12/2012	Volkert	N/A	N/A
8515736	12/2012	Duta	704/9	G06F 40/30
8761373	12/2013	Raghavan	379/111	H04M 3/4936
8838434	12/2013	Liu	704/2	G06F 40/58
8983840	12/2014	Deshmukh	704/257	G10L 15/18
9131369	12/2014	Ganong, III	N/A	H04W 12/02
9355110	12/2015	Chi	N/A	N/A
9405832	12/2015	Edwards	N/A	G06F 16/951
9484021	12/2015	Mairesse	N/A	G10L 15/08
9514747	12/2015	Bisani et al.	N/A	N/A
9558735	12/2016	Roberts et al.	N/A	N/A
9607617	12/2016	Hebert	N/A	G10L 15/18
9619459	12/2016	Hebert	N/A	G06F 16/90332
9966065	12/2017	Gruber et al.	N/A	N/A
10018977	12/2017	Cipollo et al.	N/A	N/A
10388277	12/2018	Ghosh	N/A	G10L 15/1822
10515637	12/2018	Devries et al.	N/A	N/A
10521189	12/2018	Ryabov et al.	N/A	N/A
10629186	12/2019	Slifka	N/A	N/A
2004/0049389	12/2003	Marko et al.	N/A	N/A
2006/0025995	12/2005	Erhart	704/239	G10L 15/1822
2006/0149544	12/2005	Hakkani-Tur	704/236	G10L 15/22
2008/0059188	12/2007	Konopka et al.	N/A	N/A
2008/0189390	12/2007	Heller et al.	N/A	N/A
2009/0030697	12/2008	Cerra et al.	N/A	N/A
2009/0030698	12/2008	Cerra et al.	N/A	N/A

2010/0088100	12/2009	Lindhahl	N/A	N/A
2010/0169075	12/2009	Raffa et al.	N/A	N/A
2010/0268536	12/2009	Suendermann	704/243	G10L 15/063
2010/0332234	12/2009	Agapi	704/275	G10L 15/22
2011/0038613	12/2010	Buchheit	N/A	N/A
2011/0066634	12/2010	Phillips et al.	N/A	N/A
2013/0007208	12/2012	Tsui et al.	N/A	N/A
2013/0085586	12/2012	Parekh	N/A	N/A
2013/0151250	12/2012	VanBlon	704/235	G10L 15/32
2013/0159000	12/2012	Ju	704/254	G10L 15/1822
2013/0173765	12/2012	Korbecki	N/A	N/A
2013/0326353	12/2012	Singhal	715/728	G06F 3/167
2014/0019573	12/2013	Swift	N/A	N/A
2014/0039899	12/2013	Cross, Jr.	704/275	G10L 15/193
2014/0046876	12/2013	Zhang et al.	N/A	N/A
2014/0058732	12/2013	Labsky	704/254	G10L 15/30
2014/0207442	12/2013	Ganong, III	704/201	H04W 12/02
2014/0274203	12/2013	Ganong, III et al.	N/A	N/A
2015/0012271	12/2014	Peng	704/235	G10L 15/08
2015/0120288	12/2014	Thomson	704/231	G10L 15/22
2015/0120296	12/2014	Stern et al.	N/A	N/A
2015/0134334	12/2014	Sachidanandam et al.	N/A	N/A
2015/0199967	12/2014	Reddy et al.	N/A	N/A
2015/0279352	12/2014	Willett	704/231	G10L 15/30
2015/0348548	12/2014	Piernot et al.	N/A	N/A
2016/0012819	12/2015	Willett	704/233	G10L 15/20
2016/0098998	12/2015	Wang et al.	N/A	N/A
2016/0103652	12/2015	Kuniansky	N/A	N/A
2016/0342383	12/2015	Parekh	N/A	N/A
2016/0378747	12/2015	Orr et al.	N/A	N/A
2016/0379626	12/2015	Deisher	704/232	G10L 15/30
2017/0097618	12/2016	Cipollo et al.	N/A	N/A
2017/0177716	12/2016	Perez	N/A	G06F 16/3344
2017/0213546	12/2016	Gilbert	N/A	G10L 15/1815
2017/0236512	12/2016	Williams	381/79	G10L 15/22
2017/0263253	12/2016	Thomson et al.	N/A	N/A
2017/0278511	12/2016	Willett	N/A	G10L 25/84
2017/0278514	12/2016	Mathias	N/A	G10L 15/26
2017/0294184	12/2016	Bradley	N/A	G10L 15/04
2018/0018959	12/2017	Des Jardins et al.	N/A	N/A
2018/0060326	12/2017	Kuo	N/A	G06Q 50/01
2018/0061403	12/2017	Devaraj	N/A	G10L 15/22
2018/0061404	12/2017	Devaraj	N/A	H04L 51/10
2018/0197545	12/2017	Willett	N/A	G10L 15/30
2018/0211663	12/2017	Shin	N/A	G10L 15/1815
2018/0211668	12/2017	Willett	N/A	G10L 15/30
2018/0233141	12/2017	Solomon et al.	N/A	N/A
2018/0247065	12/2017	Rhee	N/A	G07C 9/257
2018/0268818	12/2017	Schoenmackers et al.	N/A	N/A
2018/0294001	12/2017	Kayama	N/A	N/A

2018/0314689	12/2017	Wang	N/A	G06F 40/58
2018/0330728	12/2017	Gruenstein	N/A	G10L 17/005
2019/0027147	12/2018	Diamant et al.	N/A	N/A
2019/0042539	12/2018	Kalsi et al.	N/A	N/A
2019/0043509	12/2018	Suppappola	N/A	G10L 17/00
2019/0043529	12/2018	Muchlinski	N/A	G10L 25/84
2019/0057693	12/2018	Fry	N/A	G10L 15/22
2019/0103101	12/2018	Danila	N/A	G10L 25/51
2019/0295552	12/2018	Pasko	N/A	G10L 15/30
2019/0371307	12/2018	Zhao et al.	N/A	N/A
2019/0392836	12/2018	Kang et al.	N/A	N/A
2020/0051547	12/2019	Shanmugam et al.	N/A	N/A
2020/0410996	12/2019	Brandel et al.	N/A	N/A

OTHER PUBLICATIONS

Office Action for U.S. Appl. No. 15/934,761, mailed on Feb. 6, 2020, Pasko, “Speech Interface Device With Caching Component”, 11 Pages. cited by applicant

Office Action for U.S. Appl. No. 17/018,279, mailed 12/20/202, Pasko, “Speech Interface Device With Caching Component ”, 24 pages. cited by applicant

Office Action for U.S. Appl. No. 17/018,279, mailed Aug. 18, 2021, Pasko, “Speech Interface Device With Caching Component ”, 20 pages. cited by applicant

Office Action for U.S. Appl. No. 17/902,519, mailed on May 22, 2023, Inventor #1Stanislaw Ignacy Pasko, “Speech Interface Device With Caching Component ,” 8 pages. cited by applicant

Primary Examiner: Islam; Mohammad K

Attorney, Agent or Firm: Eversheds Sutherland (US) LLP

Background/Summary

CROSS REFERENCE TO RELATED APPLICATION (1) This application is a continuation of, commonly owned U.S. patent application Ser. No. 17/902,519, filed on Sep. 2, 2022, which is a continuation of U.S. patent application Ser. No. 17/018,279, filed on Sep. 11, 2020, which is a continuation of U.S. patent application Ser. No. 15/934,761, filed Mar. 23, 2018, and entitled “SPEECH INTERFACE DEVICE WITH CACHING COMPONENT,” now known as U.S. Pat. No. 10,777,203, issued on Sep. 15, 2020, the entirety of which is herein incorporated by reference.

BACKGROUND

(1) Homes and other user premises are increasingly equipped with always-on Internet or “cloud” connectivity. The constant, or nearly constant, availability of wide area network communications, in combination with increasing capabilities of computing devices—including hands-free, speech interface devices—have created a number of new possibilities for services that use voice assistant technology with in-home connected devices. For example, various cloud-based services (e.g., music streaming, smart home control, etc.) may be accessible to users through convenient, hands-free interaction with their in-home speech interface devices.

(2) Provided herein are technical solutions to improve and enhance these and other systems.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The detailed description is described with reference to the accompanying drawings. In the figures, the left-most digit(s) of a reference number identifies the figure in which the reference number first appears. The use of the same reference numbers in different figures indicates similar or identical components or features.
- (2) FIG. 1 is a block diagram illustrating a system including a hybrid speech interface device that is capable of caching remote speech processing information to enhance local speech processing on the hybrid speech interface device.
- (3) FIG. 2 is a flow diagram of an example process implemented by a hybrid speech interface device for caching remote speech processing information when an automatic speech recognition (ASR)-level difference is detected between a local speech processing component and a remote speech processing component.
- (4) FIG. 3 is a schematic diagram showing example caching logic for an ASR-level difference between a local speech processing component and a remote speech processing component.
- (5) FIG. 4 is a flow diagram of an example process implemented by a hybrid speech interface device for caching remote speech processing information when ASR results match, but a natural language understanding (NLU)-level difference is detected between a local speech processing component and a remote speech processing component.
- (6) FIG. 5 is a schematic diagram showing example caching logic for a NLU-level difference between a local speech processing component and a remote speech processing component.
- (7) FIG. 6 is a flow diagram of an example process implemented by a hybrid speech interface device for using cached information for local speech processing.
- (8) FIG. 7 is a flow diagram of an example process implemented by a hybrid speech interface device for selectively cached remote directive data based on popularity of an utterance or a likelihood of detecting the utterance in the future.
- (9) FIG. 8 illustrates example components of an electronic device, such as the hybrid speech interface device of FIG. 1.

DETAILED DESCRIPTION

- (10) Described herein are, among other things, techniques, devices, and systems, including a speech interface device with a caching component that allows for storing (e.g., caching) remote speech processing information, and using the stored (e.g., cached) information locally on the speech interface device when responding to user speech in the future. The speech interface device is also configured with “hybrid” functionality, which allows the speech interface device to process user speech locally, and to decide whether to respond to user speech using response data from a remote speech processing system, or response data from a local speech processing component. The caching component, in combination with this hybrid functionality, allows the speech interface device to respond to user speech, even in instances when a remote system—which, when available, can be used for processing user speech remotely—is, for example, unavailable to, slower than (with network-related latency factored in), or otherwise less preferred than the speech interface device.
- (11) For instance, a speech interface device may reside within an environment (e.g., in a home, automobile, office, hotel, etc.), perhaps along with one or more additional devices (including, for example, other speech interface devices, one or more second devices, such as home automation devices, mobile phone, tablet, TV, wireless speaker, etc.). In a “connected” condition, the speech interface device is able to connect to a remote system over a wide area network so that the remote system can provide speech processing for responding to user speech that occurs locally in the environment. Traditionally, the response data that is received from the remote speech processing system includes directive data that tells the speech interface device how to respond to the user speech. In the disclosed embodiments, this response data from the remote speech processing system may be enhanced with additional information and sent to the speech interface device in

real-time while the remote speech processing system is assisting the speech interface device in responding to user speech. For instance, the speech interface device described herein is configured to “listen” for utterances of a user in the environment of the speech interface device, and/or receive audio data from other speech interface devices in the environment that detect such utterances. In either case, when a remote speech processing component is available to the speech interface device, the speech interface device may send the audio data to the remote speech processing component, and may receive response data therefrom. The response data received from the remote speech processing system—in addition to including remote directive data—is enhanced with a remote ASR result(s) (e.g., remote ASR data, such as text data corresponding to the audio data), and a remote NLU result(s) (e.g., remote NLU data, such as intent data and/or slot data corresponding to the remote ASR result). In some embodiments, the response data from the remote speech processing system may include one or more cacheable status indicators associated with the NLU result(s) and/or the remote directive data, which indicate whether the remote NLU result(s) and/or the remote directive data are individually cacheable. The speech interface device can then determine whether to store (e.g., cache) at least some of this information. After storing (e.g., caching) the information on the speech interface device, the speech interface device can use the stored information to respond to user speech in the future without receiving data, or by receiving less data, from the remote speech processing system.

(12) In order to determine whether to store (e.g., cache) any of the remote information, the speech interface device may leverage its “hybrid” capabilities to process user speech locally on the speech interface device, and may then compare the local ASR and NLU results to the remote ASR and NLU results it receives from the remote speech processing system. If the speech interface device detects an ASR-level difference based on this comparison, the local speech processing component of the speech interface device may perform NLU processing on the remote ASR result in an attempt to generate the same NLU result locally. In this scenario, if a second local NLU result (based on the remote ASR result) matches the remote NLU result, the speech interface device may store the remote ASR result in local memory, along with association data that associates the remote ASR result with the first (original) local ASR result. At some later point in time, if the user utters the same expression, the local speech processing component may retrieve the remote ASR result from local memory, and may perform NLU processing on the retrieved remote ASR result using its local speech processing component in order to respond to the user speech. Not only does this allow the speech interface device to respond to user speech in instances when the Internet is down, and/or when a remote system cannot otherwise be reached, but it allows the speech interface device to do so with improved accuracy and/or enhanced functionality, as compared to an accuracy that is achievable when relying exclusively on a local speech processing component to generate ASR and NLU results. To illustrate, consider a case where a storm causes the Internet to go down, and the remote system is unavailable to the speech interface device. In this scenario, if a user utters the expression “Turn off the cellar light,” the local ASR component may generate “Turn off the seller light” as the local ASR result, where the word “cellar” is misspelled. However, assuming a remote ASR result is already stored in local memory and associated with this local ASR result, the speech interface device may nevertheless respond correctly to the utterance “Turn off the cellar light” by processing the retrieved remote ASR result through the local NLU component, causing the speech interface device to send the appropriate “turn off” command to a light in the cellar of the user's house.

(13) In some cases, a NLU-level difference may be detected despite local and remote ASR results that match. When the local and remote ASR results match, but the NLU results differ, the speech interface device may store the remote NLU result that was received from the remote speech processing system in local memory (assuming the remote NLU result is cacheable), along with association data that associates the remote NLU result with the first (original) local ASR result. At some later point in time, if the user utters the same expression, the local speech processing

component may retrieve the remote NLU result from local memory, and perform an action based at least in part on the retrieved remote NLU result. To illustrate, consider a case where a user is driving in an automobile that is also a speech interface device, and that the user is driving through a remote geographical area without connectivity to a remote speech processing system. In this scenario, if the user utters the expression “Take me home now,” the local speech processing component may not be able to resolve the intent of this utterance (e.g., a local NLU failure) due to the word “now” being incorrectly interpreted as part of the destination location. However, assuming a remote NLU result (e.g., intent data) is stored in local memory and associated with the ASR text: “Take me home now,” the speech interface device may nevertheless respond correctly to this utterance by executing the retrieved remote NLU result in order to navigate the user to his/her home address.

(14) FIG. 1 is a block diagram illustrating a system **100** including a hybrid speech interface device **102** (often shortened herein to “speech interface device **102**”) that is capable of storing (e.g., caching) remote speech processing information received from a remote system **104**. This caching allows for enhancing local speech processing on the speech interface device **102**, which is enabled by the speech interface device's **102** “hybrid” capabilities. The speech interface device **102** may be located within an environment to provide various capabilities to a user **106**, when the user **106** is also in the environment. The environment in which the speech interface device **102** is located may be a home or other premises, an automobile, or any similar environment. Such an environment may include other devices including additional speech interface devices, such as the speech interface device **108**, and/or second devices (e.g., Internet of Things (IoT) devices and/or smart home devices like thermostats, lights, refrigerators, ovens, etc.) that may be controllable by speech interface devices, such as the speech interface device **102**. When acting as a hub, the speech interface device **102** may be configured to connect a plurality of devices in an environment and control communications among them, thereby serving as a place of convergence where data arrives from one or more devices, and from which data is sent to one or more devices.

(15) In general, the speech interface device **102** may be capable of capturing utterances with a microphone(s) **110**, and responding in various ways, such as by outputting content via an output device(s) **112**, which may be a speaker(s), a display(s), or any other suitable output device **112**. In addition, the speech interface device **102** may be configured to respond to user speech by controlling second devices that are collocated in the environment with the speech interface device **102**, such as by sending a command to a second device via a communications interface **114** (e.g., a short range radio), the command instructing an operation to be performed at the second device (e.g., to turn on a light in the environment). FIG. 1 also shows that, in addition to using the microphone(s) **110** to capture utterances as audio data **116**, the speech interface device **102** may additionally, or alternatively, receive audio data **116** (e.g., via the communications interface **114**) from another speech interface device **108** in the environment, such as when the other speech interface device **108** captures an utterance from the user **106** and sends the audio data **116** to the speech interface device **102**. This may occur in situations where the other speech interface device **108** would like to leverage the “hybrid” capabilities of the speech interface device **102**.

(16) As mentioned, under normal conditions, the speech interface device **102** may operate in conjunction with and/or under the control of a remote, network-based or network-accessible control system **104** (abbreviated to “remote system” **104** in FIG. 1). The remote system **104** may, in some instances be part of a network-accessible computing platform that is maintained and accessible via a wide area network **118**. Network-accessible computing platforms such as this may be referred to using terms such as “on-demand computing”, “software as a service (SaaS)”, “platform computing”, “network-accessible platform”, “cloud services”, “data centers”, and so forth. The remote system **104** may be configured to provide particular functionality to large numbers of local (e.g., in-home, in-car, etc.) devices of different users.

(17) In some embodiments, the remote system **104** may be configured to receive audio data **116**

from the speech interface device **102**, to recognize speech in the received audio data **116** using a remote speech processing system **120**, and to perform functions in response to the recognized speech. In some embodiments, these functions involve sending directives, from the remote system **104**, to the speech interface device **102** to cause the speech interface device **102** to perform an action, such as output an audible response to the user speech via a speaker(s) (i.e., an output device(s) **112**), and/or control second devices in the environment by sending a control command via the communications interface **114**. Thus, under normal conditions, when the speech interface device **102** is able to communicate with the remote system **104** over a wide area network **118** (e.g., the Internet), some or all of the functions capable of being performed by the remote system **104** may be performed by sending a directive(s) over the wide area network **118** to the speech interface device **102**, which, in turn, may process the directive(s) for performing an action(s). For example, the remote system **104**, using a remote directive that is included in a remote response **121** (e.g., response data), may instruct the speech interface device **102** to output an audible response (e.g., using text-to-speech (TTS)) to a user's question, to output content (e.g., music) via a speaker of the speech interface device **102**, and/or to turn on a light in the environment. It is to be appreciated that the remote system **104** may be configured to provide other functions in addition to those discussed herein, such as, without limitation, providing step-by-step directions for navigating from an origin to a destination location, conducting an electronic commerce transaction on behalf of the user **106** as part of a shopping function, establishing a communication session between the user **106** and another user, and so on.

(18) A home automation example is shown in FIG. **1** by the user **106** uttering the expression “Turn off the cellar light.” Whether this utterance is captured by the microphone(s) **110** of the speech interface device **102** or captured by another speech interface device **108** in the environment, the audio data **116** representing this user's speech is ultimately received by a speech interaction manager (SIM) **122** of a voice services component **124** executing on the speech interface device **102**. The SIM **122** may manage received audio data **116** by processing utterances as events, and the SIM **122** may also manage the processing of directives that are used to respond to the user speech (e.g., by controlling the action(s) of the speech interface device **102**). The SIM **122** may include one or more client applications **126** or skills for performing various functions at the speech interface device **102**.

(19) A hybrid request selector **128** (or, hybrid request selector component **128**) is shown as a component of the speech interface device **102** in FIG. **1**. At least part of (e.g., subcomponents of) the hybrid request selector **128** can be implemented as a layer within the voice services component **124** that is located between the SIM **122** and a speech communication library (SCL) **132**, and may be configured to proxy traffic to/from the remote system **104**. For example, the hybrid request selector **128** may be configured to pass messages between the SIM **122** and the SCL **132** (such as by passing events and directives there between), and to send messages to/from subcomponents of the hybrid request selector **128** that are configured to select a response (e.g., either local or remote response data) to use in responding to user speech. In this manner, the hybrid request selector **128** can “snoop” on communication between the SIM **122** and the remote system **104**. For instance, information (e.g., a remote directive, a capability indicator, etc.) that is contained in the remote response **121** can be sent to the hybrid request selector **128**. The hybrid request selector **128** may also be configured to allow audio data **116** received from the SIM **122** to pass through to the remote speech processing system **120** (via the SCL **132**) while also receiving (e.g., intercepting) this audio data **116** for input to a local speech processing component **140**.

(20) The local speech processing component **140** is configured to process audio data **116** representing user speech. In some embodiments, the hybrid request selector **128** may further control the execution of the local speech processing component **140**, such as by sending “execute” and “terminate” events/instructions to the local speech processing component **140**. An “execute” event may instruct the local speech processing component **140** to continue any suspended

execution based on audio data **116** (e.g., by instructing the local speech processing component **140** to execute on a previously-determined intent in order to generate a directive). Meanwhile, a “terminate” event may instruct the local speech processing component **140** to terminate further execution based on the audio data **116**, such as when the speech interface device **102** receives directive data from the remote system **104** and chooses to use that remotely-generated directive data.

(21) In the example of FIG. **1**, where the user **106** utters the expression “Turn off the cellar light,” the audio data **116** is received by the SIM **122** of the voice services component **124**. The SIM **122** may send the audio data **116** through the hybrid request selector **128**, which allows the audio data **116** to pass through to the SCL **132**, and the SCL **132**, in turn, sends the audio data **116** over the wide area network **118** to the remote system **104** for processing speech remotely. The wide area network **118** is representative of any type of public or private, wide-area network, such as the Internet, which extends beyond the environment of the speech interface device **102**. Thus, the wide area network **118** may represent and/or include, without limitation, data and/or voice networks, a wired infrastructure (e.g., coaxial cable, fiber optic cable, etc.), a wireless infrastructure (e.g., radio frequencies (RF), cellular, satellite, etc.), and/or other connection technologies.

(22) The hybrid request selector **128** may also receive the audio data **116** and input the audio data **116** to the local speech processing component **140**. At this point, the hybrid request selector **128** may wait for response data from either or both of the remote system **104** or the local speech processing component **140**.

(23) The local speech processing component **140** is configured to receive the audio data **116** from the hybrid request selector **128** as input, to recognize speech in the audio data **116**, to determine user intent from the recognized speech, and to determine how to act on the user intent by generating directive data. In some cases, a directive may include a description of the intent (e.g., an intent to turn off {device A}). In some cases, a directive may include (e.g., encode) an identifier of a second device, such as the cellar light, and an operation to be performed at the second device.

(24) In some embodiments, the speech interface device **102** may be configured to compare the audio data **116** to stored models used to detect a wakeword that indicates to the speech interface device **102** that the audio data **116** is to be processed for determining an intent (a local NLU result). In some embodiments, the hybrid request selector **128** may send the audio data **116** to the local speech processing component **140** to have the local speech processing component **140** determine whether a wakeword is detected in the audio data **116**, and, if a wakeword is detected, this indication may be provided to the hybrid request selector **128** so that the hybrid request selector **128** can proceed with routing the audio data **116** to the remote speech processing system **120** while the local speech processing component **140** continues processing the audio data **116** in parallel. If the local speech processing component **140** does not detect a wakeword in the audio data **116**, this indication may be provided to the hybrid request selector **128** so that the hybrid request selector **128** can refrain from sending the audio data **116** to the remote speech processing system **120**, and the local speech processing component **140** may halt further operations after determining that the audio data **116** does not include the wakeword. The audio data **116** can be discarded in this situation.

(25) Among other logical and physical components, the local speech processing component **140** may include an automatic speech recognition (ASR) component **142** that is configured to perform ASR processing on the audio data **116** to convert the audio data **116** into ASR text data (referred to herein as an “ASR result”, or “ASR data”). ASR transcribes audio data into text data representing the words of the user speech contained in the audio data **116**. A spoken utterance in the audio data can be input to the ASR component **142**, which then interprets the utterance based on the similarity between the utterance and pre-established language models available to the local speech processing component **140**. For example, the ASR component **142** may compare the input audio data **116** with models for sounds (e.g., subword units or phonemes) and sequences of sounds to identify words

that match the sequence of sounds spoken in the utterance of the audio data **116**. In some embodiments, the ASR component **142** outputs the most likely text recognized in the audio data **116**, or multiple hypotheses in the form of a lattice or an N-best list with individual hypotheses corresponding to confidence scores or other scores (such as probability scores, etc.). In some embodiments, the ASR component **142** is customized to the user **106** (or multiple users) who created a user account to which the speech interface device **102** is registered. For instance, the language models (and other data) used by the ASR component **142** may be based on known information (e.g., preferences) of the user **106**, and/or on a history of previous interactions with the user **106**.

(26) The local speech processing component **140** may also include a natural language understanding (NLU) component **144** that performs NLU processing on the generated ASR text data to determine intent data and/or slot data (referred to herein as a “NLU result”, or “NLU data”) so that directives may be determined based on the intent data and/or the slot data. Generally, the NLU component **144** takes textual input (such as processed from the ASR component **142**) and attempts to make a semantic interpretation of the ASR text data. That is, the NLU component **144** determines the meaning behind the ASR text data based on the individual words, and then the NLU component **144** can implement that meaning. The NLU component **144** interprets a text string to derive an intent or a desired action or operation from the user **106**. This may include deriving pertinent pieces of information in the text that allow the NLU component **114** to identify a second device in the environment, if the user, for example, intends to control a second device (e.g., a light in the user's **106** house, as is the case in the example of FIG. 1). The local speech processing component **140** may also provide a dialog management function to engage in speech dialogue with the user **106** to determine (e.g., clarify) user intents by asking the user **106** for information using speech prompts. In some embodiments, the NLU component **144** is customized to the user **106** (or multiple users) who created a user account to which the speech interface device **102** is registered. For instance, data used by the NLU component **144** to understand the meaning of ASR text may be based on known information (e.g., preferences) of the user **106**, and/or on a history of previous interactions with the user **106**.

(27) The local speech processing component **140** may also include, or be configured to use, one or more installed speechlets **146**. Speechlets **146** may represent domains that are used in order to determine how to act on an utterance in a particular way, such as by outputting a directive that corresponds to the determined intent, and which can be processed to implement the desired operation. Accordingly, the term “speechlet” may be used interchangeably with the term “domain” or “domain implementation.” The speechlets **146** installed on the speech interface device **102** may include, without limitation, a music speechlet **146** (or music domain) to act on utterances with intents to play music on a device, such as via a speaker(s) of the speech interface device **102**, a navigation speechlet **146** (or a navigation domain) to act on utterances with intents to get directions to a point of interest with a known address, a shopping speechlet **146** (or shopping domain) to act on utterances with intents to buy an item from an electronic marketplace, and/or a device control speechlet **146** (or device control domain) to act on utterances with intents to control a second device(s) in the environment.

(28) In order to generate a particular interpreted response, the NLU component **144** may apply grammar models and lexical information associated with the respective domains or speechlets **146** to recognize one or more entities in the text of the query. In this manner the NLU component **144** may identify “slots” (i.e., particular words in query text) that may be needed for later command processing. Depending on the complexity of the NLU component **144**, it may also label each slot with a type of varying levels of specificity (such as noun, place, city, artist name, song name, device name, or the like). Each grammar model used by the NLU component **144** may include the names of entities (i.e., nouns) commonly found in speech about the particular domain (i.e., generic terms), whereas the lexical information (e.g., from a gazetteer) is personalized to the user(s) and/or

the device. For instance, a grammar model associated with the navigation domain may include a database of words commonly used when people discuss navigation.

(29) Accordingly, the intents identified by the NLU component **144** may be linked to domain-specific grammar frameworks with “slots” or “fields” to be filled (e.g., resolved). Each slot/field corresponds to a portion of the query text that the system believes corresponds to a named entity. For example, if “play music” is an identified intent, a grammar framework(s) may correspond to sentence structures such as “Play {Artist Name},” “Play {Album Name},” “Play {Song name},” “Play {Song name} by {Artist Name},” etc. However, to make slot resolution more flexible, these frameworks would ordinarily not be structured as sentences, but rather based on associating slots with grammatical tags.

(30) For example, the NLU component **144** may parse the query to identify words as subject, object, verb, preposition, etc., based on grammar rules and/or models, prior to recognizing named entities. The identified verb may be used by the NLU component **144** to identify an intent, which is then used to identify frameworks. A framework for an intent of “play” may specify a list of slots/fields applicable to play the identified “object” and any object modifier (e.g., a prepositional phrase), such as {Artist Name}, {Album Name}, {Song name}, etc. The NLU component **144** may then search the corresponding fields in the domain-specific and personalized lexicon(s), attempting to match words and phrases in the query tagged as a grammatical object or object modifier with those identified in the database(s).

(31) After determining a local NLU result, or failing to do so, the local speech processing component **140** may send response data to the hybrid request selector **128**, such as a “ReadyToExecute” response, which indicates that the local speech processing component **140** has recognized an intent, or is ready to communicate failure (if the local speech processing component **140** could not recognize an intent via the NLU component **144**). The hybrid request selector **128** may then determine whether to use directive data from the local speech processing component **140** to respond to the user speech, or whether to use directive data received from the remote system **104**, assuming a remote response **121** is even received (e.g., when the speech interface device **102** is able to access the remote speech processing system **120** over the wide area network **118**).

(32) The local speech processing component **140** may also send local ASR results and local NLU results to a caching component **148**. Additionally, the caching component **148** may receive information included in the remote response **121** via the hybrid request selector **128**, at least in instances when a remote response **121** is received. Thus, when the speech interface device **102** is in a connected condition, the caching component **148** may possess the local ASR result(s) and the local NLU result(s), as well as information extracted from the remote response **121**, which may include at least: one or more remote ASR results, one or more remote NLU results, and one or more directives, along with supplemental information, such as cacheable status indicators, skill ID(s), static response indicators, as will be described in more detail below.

(33) With both local and remote information in its possession, the caching component **148** is configured to determine whether to cache some or all of the cacheable information extracted from the remote response **121** (e.g., response data). Caching logic may be based at least in part on detecting ASR-level differences, and/or detecting NLU-level differences, as well as other factors, as will be described in more detail below. When the caching component **148** determines to do so, it may cache remote speech processing information in local memory of the speech interface device **102** as cache entries **150**. These cache entries **150** may include, without limitation, remote ASR results associated with local ASR results, and/or remote NLU results associated with local ASR results, and/or remote directives (including, e.g., static responses) associated with local ASR results, which enables the caching component **148** to retrieve the remote ASR results and/or the remote NLU results (and possibly the remote directive data) in instances where the hybrid request selector **128** determines to use the local speech processing component **140** without reliance on the remote speech processing system **120**. For instance, if the cache entries **150** include a correction

(e.g., a remote ASR result associated with the local ASR result, and/or a local NLU result associated with the local ASR result), these “corrections” can be retrieved and used in local speech processing, in lieu of generating a local ASR result and/or a local NLU result using the local speech processing component **140**.

(34) In some embodiments, a static skill handler **152** (or, static skill handler component **152**) may be configured to output static responses that are also cached local memory of the speech interfaced device (e.g., within the cache entries **150**). These static responses may include, without limitation, recordings of short songs (e.g., a recording of the Happy Birthday Song sung by a voice assistant personality), or a canned response to a user's question, as spoken by the voice assistant personality. Thus, instead of generating directive data locally for purposes of responding to user speech, the caching component **148**, or the static skill handler **152** via the caching component **152**, may be configured to retrieve static responses that are output via a static skill handler **152**. In this scenario, the SIM **122** may not process directive data, as it typically does, to cause an action to be performed by the speech interface device **102**.

(35) The caching-related features described herein allow the speech interface device **102** equipped with local speech processing component **140** to expand and enhance its own, local (e.g., offline) speech processing capabilities by leveraging results produced by the remote speech processing system **120** and received while the speech interface device **102** is in a connected condition. It is a cheap way to implement at least some of the advanced intelligence that may be available to the remote speech processing system **120**, and which is not typically available to a resource-constrained device, such as the speech interface device **102**.

(36) In some embodiments, the speech interface device **102** may share cached information with other speech interface devices in its environment, such as the speech interface device **108** of FIG. **1**. For instance, the speech interface device **102** may send a remote ASR result(s) and/or a remote NLU result(s), and/or remote directive data—along with association data that associates such remote speech processing information with a locally-generated ASR result—to the speech interface device **108**. Such transmission of cached information may be sent via a short range wireless protocol (e.g., Bluetooth®), via a local area network (LAN), etc., so that other speech interface devices can benefit from another's cached information. In some embodiments, a speech interface device **102** may selectively share a subset of its cached information with other speech interface devices, such as selecting cached information to share based on domain (e.g., music domain, shopping domain, etc.), based on device capabilities, etc.

(37) The processes described herein are illustrated as a collection of blocks in a logical flow graph, which represent a sequence of operations that can be implemented in hardware, software, or a combination thereof. In the context of software, the blocks represent computer-executable instructions that, when executed by one or more processors, perform the recited operations. Generally, computer-executable instructions include routines, programs, objects, components, data structures, and the like that perform particular functions or implement particular abstract data types. The order in which the operations are described is not intended to be construed as a limitation, and any number of the described blocks can be combined in any order and/or in parallel to implement the processes.

(38) FIG. **2** is a flow diagram of an example process **200** implemented by a hybrid speech interface device **102** for caching remote speech processing information when an ASR-level difference is detected between a local speech processing component **140** and a remote speech processing system **120**. For discussion purposes, the process **200** is described with reference to FIGS. **1** and **3**.

(39) At **202**, a local speech processing component **140** executing on a speech interface device **102** may perform ASR processing on audio data **116** that represents user speech to generate a first local ASR result. Turning briefly to FIG. **3**, an example utterance **300** that might be detected by a speech interface device, such as the speech interface device **102**, is shown as the utterance “Turn off the cellar light.” FIG. **3** then shows how the operation(s) at block **202** may be performed to generate a

first local ASR result **302** (e.g., “Turn off the seller light”). That is, the first local ASR result **302** may be generated by the local ASR component **142** performing ASR processing on the audio data **116** that represents the user speech “Turn off the cellar light.” In this case, the local ASR component **142** generated a first local ASR result **302** with the word “cellar” misspelled as “seller.” (40) At **204** of the process **200**, the local speech processing component **140** may perform NLU processing on the first local ASR result **302** to generate a first local NLU result. In FIG. 3, an example first local NLU result **304** is shown as a failure to recognize an intent. It is to be appreciated that, if the local NLU component **144** is able to resolve the first local ASR result **302** into an intent, the first local NLU result **304** may be a recognized intent, rather than a failure to recognize an intent. In the illustrative example, the NLU component **144** was unable to resolve the first local ASR result **302** into an intent, so the first local NLU result **304** represents a failure to recognize an intent.

(41) At **206** of the process **200**, the speech interface device **102** may send the audio data **116** to a remote speech processing system **120** executing on a remote system **104**. The audio data **116** may be send over a wide area network **118** at block **206**.

(42) At **208**, the speech interface device **102** may receive response data from the remote system **104** (e.g., the remote response **121** of FIG. 1). The remote response **121** received from the remote system **104** may include: a remote ASR result(s) corresponding to the audio data **116**, a remote NLU result(s) corresponding to the remote ASR result(s), and remote directive data. This remote response **121** (e.g., response data) can be received by the hybrid request selector **128** and forwarded to the caching component **148** at block **208**. FIG. 3 shows examples of a remote ASR result **306**, a remote NLU result **308**, and remote directive data **310**. For example, the remote ASR result **306** might be “Turn off the cellar light” based on the utterance **300**, the remote NLU result **308** might be “Turn off {device A}” (where “device A”, in this case, may be the text of a resolved slot for a “turn off” intent), and the remote directive data **310** may be generated by the remote speech processing system **120** executing on the remote NLU result **308** (i.e., an intent determined by the remote speech processing system **120**). It is to be appreciated that the remote system **104** may send multiple directives related to the processing of the remote NLU result **308**. For simplicity, the example of FIG. 3 assumes a single directive, which causes the speech interface device **102** to send a “turn off” command to a light in the user's house.

(43) In some embodiments, the remote ASR result(s) **306** is received as an ASR recognition event, and may include an N-best list of remote ASR results **306**. Remote ASR results **306** may be cacheable by default, and may, therefore, omit a cacheable status indicator. In some embodiments, the remote NLU result(s) **308** is received as a NLU recognition event, and may include an N-best list of remote NLU results **308** (e.g., intent resolutions). FIG. 3 shows that the remote NLU result(s) **308** may include, or be associated with, a cacheable status indicator (e.g., “C”) to indicate whether the remote NLU result(s) **308** is cacheable; in this case, “C” is to be interpreted as a cacheable status indicator that indicates something is cacheable. In some embodiments, where a NLU result **308** includes resolved (e.g., filled) slots, a cacheable status indicator (e.g., “C”) may be associated with each slot to indicate whether the slot itself (and/or the resolved text of the slot) is cacheable. In this scenario, a NLU result **308** may be cacheable if all of its slots are cacheable. A NLU result **308** may be cacheable if the NLU result **308** does not contain slots that may change their interpretation over time. For example, the utterance “what's the weather in Houston” may be resolved into a cacheable remote intent (remote NLU result **308**). Remote NLU results **308** for the utterances “what's the time” or “tell me a joke” may be cacheable as well, because their interpretations don't change over time. It is to be appreciated that, although the letter “C” is used to denote a cacheable status indicator herein, a cacheable status indicator may be provided in any suitable form, such as a bit or Boolean value that can be set to “true” to indicate that information is cacheable, or “false” to indicate that information is not cacheable.

(44) The remote directive data **310** may also be associated with a cacheable status indicator (e.g.,

“C”), which indicates whether the remote directive data **310** is cacheable. Remote directive data **310** may be cacheable if the remote directive data **310** does not change over time. In the example where the user **106** utters the expression “what's the weather in Houston”, the remote directive data **310** for that utterance is not cacheable because the weather changes over time, and, therefore, the remote directive data **310** for outputting an appropriate response to the user **106** from the speech interface device **102** will change over time. In the example shown in FIG. 3, the remote directive data **310** (for turning off a light in the cellar) may be cacheable. In some embodiments, the remote directive data **310** may also include a skill identifier (ID) that identifies a remote skill (e.g., an application) that executed on the remote NLU result **308**. In some embodiments, the remote directive data **310** may also include a static response indicator to indicate when a “response” (which is to be output from an output device **112** of the speech interface device **102**) is static, and, as a result, can be cached for playback on the speech interface device **102**. An example of a static response may be a recording of content (e.g., audio and/or video content), such as the Happy Birthday song, which can be cached and subsequently retrieved for playback locally on the speech interface device **102**, if needed. Static responses may be retrieved from local memory of the speech interface device **102** and executed by the static skill handler **152**.

(45) After block **208** of the process **200**, the caching component **148** has access to, among other things, the first local ASR result **302** generated by the local speech processing component **140**, and the remote ASR result **306** that was extracted from the remote response **121** received at block **208**. It is to be appreciated that the caching component **148**, at this point in time, may have access to 1 or more local ASR results **302**, 0 or more local intents that are recognized by the local NLU component **144**, 1 or more remote ASR results **306**, and 0 or more remote intents that are recognized by the remote speech processing system **120**. This is because performing NLU processing sometimes results in a failure to recognize an intent, and the NLU results **304/308** may, therefore, represent either a failure to recognize an intent, or a recognized intent.

(46) At **210** of the process **200**, the caching component **148** may compare the first local ASR result **302** to the remote ASR result **306**. This comparison at block **210** may include selecting and comparing a top ranked (e.g., 1-best) local ASR result **302** to a top ranked (e.g., 1-best) remote ASR result **306**.

(47) At **212**, the caching component **148** may determine whether, based on the comparison at block **210**, the first local ASR result **302** matches the remote ASR result **306**. It is to be appreciated that a determination as to whether first ASR text data matches second ASR text data may include determining an approximate text match between the first ASR text data and the second ASR text data. Approximate text matching may include generating a match score and determining whether the match score satisfies (e.g., meets or exceeds) a threshold score. For instance, there may be differences between first ASR text data and second ASR text data for various reasons, such as formatting differences, noise, differences in vocabulary available to the respective ASR components or the like, and a match score can be calculated based on a change metric, such as a number of character and/or word edits to change the first ASR text data into the second ASR text data, or a similar metric. An exact match would not have any character and/or word edits, and, as a result, may receive a highest possible match score. The threshold score may be set at any suitable level below the highest possible match score to allow for first and second ASR text data to be considered “matching”, despite one or more differences between the first and second ASR text data. It follows that, first ASR text data may be considered to not match, or differ from, second ASR text data if a match score does not satisfy (e.g., does not meet or exceed) a threshold score.

(48) If the local and remote ASR results match at block **212**, the process **200** may follow the “YES” route from block **212** to the off-page reference “A” (which continues in FIG. 4). If the first local ASR result **302** differs from (e.g., does not match) the remote ASR result **306** at block **212**, the process **200** may follow the “NO” route from block **212** to block **214**. The example of FIG. 3 shows a case where the first local ASR result **302** (e.g., “Turn off the seller light”) does not match

the remote ASR result (e.g., “Turn off the cellar light”) due to the difference in the words “seller” and “cellar”, and as such, the process **200** may follow the “NO” route from block **212** for the example shown in FIG. 3.

(49) At **214**, the caching component **148** may compare the first local NLU result **304** to the remote NLU result **308**. This comparison at block **214** may include selecting and comparing a top ranked (e.g., 1-best) local NLU result **304** to a top ranked (e.g., 1-best) remote NLU result **308**. Comparing the local and remote NLU results **304/308** at block **214** may include analyzing the first local NLU result **304** to determine whether the first local NLU result **304** is a failure to recognize an intent. Comparing the local and remote NLU results at block **214** may, additionally or alternatively, include comparing slot resolutions of the local and remote NLU results **304/308**.

(50) At **216** of the process **200**, the caching component **148** may determine whether, based on the comparison (e.g., analysis of the first local NLU result **304**) at block **214**, the first local NLU result **304** matches the remote NLU result **308**. It is to be appreciated that a determination as to whether first NLU intent data matches second NLU intent data may include determining an approximate text match between the first NLU intent data and the second NLU intent data. Approximate text matching may include generating a match score and determining whether the match score satisfies (e.g., meets or exceeds) a threshold score. For instance, there may be differences between first NLU intent data and second NLU intent data for various reasons, such as formatting differences, differences in intent recognition capabilities, differences in vocabulary available to the respective NLU components, or the like, and a match score can be calculated based on a change metric, such as a number of character and/or word edits to change the first NLU intent data into the second NLU intent data, or a similar metric. An exact match would not have any character and/or word edits, and, as a result, may receive a highest possible match score. The threshold score may be set at any suitable level below the highest possible match score to allow for first and second NLU intent data to be considered “matching”, despite one or more differences between the first and second NLU intent data. It follows that, first NLU intent data may be considered to not match, or differ from, second NLU intent data if a match score does not satisfy (e.g., does not meet or exceed) a threshold score. If the local and remote NLU results match each other at block **216**, the process **200** may follow the “YES” route from block **216** to block **218**.

(51) At **218**, when the first local NLU result **304** matches the remote NLU result **308**, notwithstanding a difference between the first local ASR result **302** and the remote ASR result **306**, the caching component **148** may treat this condition as a discrepancy. As shown by sub-blocks **217** and **219**, the caching component **148** may report (sub-block **217**) the discrepancy to the remote system **104** as a metric, and may upload (sub-block **219**) one or more local results (e.g., the first local ASR result **302** and/or the first local NLU result **304**) to the remote system **104**. The remote system **104** may use the reported metric and/or the uploaded local result(s) to retrain models (e.g., language models, etc.) used by local speech processing components, such as the local speech processing component **140**, in an effort to improve the local speech recognition capabilities of these local components.

(52) If, at block **216**, the first local NLU result **304** differs from (e.g., does not match) the remote NLU result **308** (which may include a determination that the first local NLU result **304** represents a failure to recognize an intent), the process **200** may follow the “NO” route from block **216** to block **220**. The example of FIG. 3 shows a case where the first local NLU result **304** (e.g., a failure to recognize an intent) does not match the remote NLU result (e.g., “Turn off {device A}”).

(53) At **220** of the process **200**, when at least one of: (i) the first local NLU result **304** represents a failure to recognize an intent, or (ii) the first local NLU result **304** does not match the remote NLU result **308**, the local speech processing component **140** (e.g., the NLU component **144**) may perform NLU processing on the remote ASR result **306** in an attempt to generate a second local NLU result **312**.

(54) At **222**, the caching component **148** may compare the second local NLU result **312** to the

remote NLU result **308**. Comparing the local and remote NLU results at block **222** may include analyzing the second local NLU result **312** to determine whether the second local NLU result **312** is a failure to recognize an intent. The comparing the local and remote NLU results at block **222** may, additionally or alternatively, include comparing slot resolutions of the local and remote NLU results.

(55) At **224**, if the caching component **148** determines that the second local NLU result **312** matches the remote NLU result **308**, the process **200** may follow the “YES” route from block **224** to block **226**. This case is shown in the example of FIG. **3**, where the local NLU component **144** resolved the remote ASR result **306** into a second local NLU result **312**: “Turn off {device A}”, which is the same as the remote NLU result **308**.

(56) At **226**, the caching component **148** may treat this condition as a correction, and may cache the remote ASR result **306** by storing, in memory of the speech interface device **102** (e.g., as a cache entry **150**): the remote ASR result **306**, and association data that associates the remote ASR result **306** with the first local ASR result **302**. As noted in FIG. **2**, this storing of remote speech processing information in the memory of the speech interface device **102** is sometimes referred to as “precaching,” because the cached information is to be retrieved and used in the future, in particular instances. This is shown in the first cache entry of the cache entries **150** depicted in FIG. **3**, where the first local ASR result **302** (e.g., “Turn off the seller light”) is associated with the remote ASR result **306** (e.g., “Turn off the cellar light”), and is indicated as a correction. In some embodiments, the correction information stored in the local memory may include a trust score associated with the first local ASR result **302**. The trust score can be used to determine when to lookup a local ASR result in the cache entries **150** (e.g., when the trust score is less than a threshold score). In some embodiments, the remote directive data **310** (when cacheable) is cached at block **226** as well, and the remote directive data **310** can be associated with the first local ASR result **302**. Caching the remote directive data **310** may include caching a static response (e.g., a recording of an audible response, such as a speech byte stream, or content that can be played back using the static skill handler **152**). As mentioned, some remote directive data **310** may not be cacheable, so storing the remote directive data **310** may include determining that a cacheable status indicator associated with the remote directive data **310** indicates that the remote directive data **310** is cacheable.

(57) If, at block **224**, the caching component **148** determines that the second local NLU result **312** does not match the remote NLU result **308** (which may include a determination that the second local NLU result **312** represents a failure to recognize an intent), the process **200** may follow the “NO” route from block **224** to block **228**, where the caching component **148** may treat this condition as a discrepancy. As shown by sub-blocks **227** and **229**, the caching component **148** may report (sub-block **227**) the discrepancy to the remote system **104** as a metric, and may upload (sub-block **229**) one or more local results (e.g., the first local ASR result **302**, the first local NLU result **304**, and/or the second local NLU result **310**) to the remote system **104**. The remote system **104** may use the reported metric and/or the uploaded local results to retrain models (e.g., language models, etc.) used by local speech processing components, such as the local speech processing component **140**, to improve the speech recognition capabilities of the local components.

(58) At **230**, following either a discrepancy condition or a correction condition, the speech interface device **102** may continue to respond to the user speech by performing an action. This action can be based on the remote directive data **310** received from the remote system **104**. Alternatively, if the hybrid request selector **128** determines to use locally-generated directive data, the action may be based on locally-generated directive data. In the example of FIG. **3**, having received a remote response **121** from the remote system **104**, the remote directive data **310** may be used to cause the speech interface device **102** to perform the action, such as outputting content (e.g., audio and/or video content) via an output device **112** of the speech interface device **102**, sending a command to a second device collocated in an environment with the speech interface device **102**, the command instructing an operation to be performed at the second device (e.g., turn on/off a light), and so on.

(59) In some embodiments, directive data that is generated by a domain/speechlet of the remote speech processing system **120** and/or the local speech processing component **140** may be formatted using Java, such as JavaScript syntax, or JavaScript-based syntax. This may include formatting the directive using JSON. In some embodiments, such a locally-generated directive may be serialized, much like how remotely-generated directives are serialized for transmission in data packets over the wide area network **118**. In other embodiments, a locally-generated directive is formatted as a programmatic application programming interface (API) call with a same logical operation as a remotely-generated directive. In other words, a locally-generated directive may mimic remotely-generated directives by using a same, or a similar, format as the remotely-generated directive.

(60) An example structure of a locally-generated directive may include a header object, an endpoint object, and a payload object. The header object may include a set of fields that provide different types of identifying information (sometimes called “properties”). Header properties may include, without limitation, a namespace (e.g., for specifying a category for the message payload, e.g., “SpeechSynthesizer,”), a name (e.g., “Speak”), a message identifier (ID) (e.g., a unique identifier for a single request or response), a correlation token (e.g., to identify a directive, and possibly events associated therewith), a payload version (e.g., version of the capability interface applied to the message), a “keys” section, etc. An endpoint object may identify a target for a directive, and possibly an origin of an event. An endpoint, in this context, can represent a physical device, a virtual device, a group of devices, a software component. The endpoint may include an authentication token to enable communication with a device(s) or component represented by the endpoint. Endpoint properties may include, without limitation, a scope (e.g., a polymorphic object that provides authorization and identifying information for a message, and may include a token that identifies a user), an endpoint ID (e.g., a unique identifier corresponding to the endpoint), and a cookie (e.g., a list of key/value pairs associated with the endpoint).

(61) FIG. 4 is a flow diagram of an example process **400** implemented by a hybrid speech interface device **102** for caching remote speech processing information when ASR results match, but an NLU-level difference is detected between a local speech processing component **140** and a remote speech processing system **120**. For discussion purposes, the process **400** is described with reference to the previous figures, and FIG. 5. Furthermore, as shown by the off-page reference “A” in FIGS. 2 and 4, the process **400** may continue from block **212** of FIG. 2, after the caching component **148** determined that a first local ASR result matched a remote ASR result.

(62) Turning briefly to FIG. 5, an example is shown where an NLU-level difference is detected notwithstanding ASR results that match between a local speech processing component **140** and a remote speech processing system **120**. For example, an utterance **500** detected by a speech interface device **102** might be “Take me home now,” such as when the user **106** is trying to invoke a navigation service by speaking to the speech interface device **102**, which may be located in the user's **106** automobile. FIG. 5 shows how a first local ASR result **502** (e.g., “Take me home now”) may be generated by the local ASR component **142** performing ASR processing on the audio data **116** that represents this user speech. In this case, the local ASR component **142** generated a first local ASR result **502** that matches a remote ASR result **506**, causing the process **200** to follow the “YES” route from block **212** to block **402** via the off-page reference “A” in FIGS. 2 and 4.

(63) At **402** of the process **400**, the caching component **148** may compare a first local NLU result to the remote NLU result. In the example of FIG. 5, the first local NLU result **504** represents a failure to recognize an intent, however, the local NLU component **144** may generate a first local NLU result **504** as a recognized intent, in some embodiments. Meanwhile, the remote NLU result **508** of FIG. 5 is a recognized intent (e.g., “Navigate to {Address A}”). Comparing the local and remote NLU results at block **402** may include analyzing the first local NLU result **504** to determine whether the first local NLU result **504** is a failure to recognize an intent. Comparing the local and remote NLU results at block **402** may, additionally or alternatively, include comparing slot resolutions of the local and remote NLU results.

(64) At **404**, the caching component **148** may determine whether the first local NLU result **504** matches the remote NLU result **508**. If the NLU results match at block **404**, the process **400** may follow the “YES” route from block **404** to block **406** where the caching component **148** may treat this condition as a validation. As shown by sub-block **407**, the caching component **148** may generate and store a trust score (sub-block **407**) associated with either or both of the first local ASR result **502** and/or the first local NLU result **504**. A high trust score (e.g., a trust score that meets or exceeds a threshold trust score), for example, can be used in the future to refrain from looking up cache entries **150** the next time the same utterance is detected, because the local speech processing component **140** may respond to the user speech correctly without reliance on any remote results.

(65) If, at block **404**, at least one of: (i) the first local NLU result **502** represents a failure to recognize an intent, or (ii) the first local NLU result **502** does not match the remote NLU result **508** (which may include different slot resolutions), the process **400** may follow the “NO” route from block **404** to block **408**.

(66) At **408** of the process **400**, the caching component **148** may determine whether the remote NLU result **508** is executable using a local speechlet **146** on the speech interface device **102**. If the remote NLU result **508** is not executable locally on the speech interface device **102** using a local speechlet **146**, the process **400** may follow the “NO” route from block **408** to block **410** where the caching component **148** may treat this condition as a discrepancy. A NLU discrepancy can be handled in various ways. In some embodiments, as shown by sub-block **409**, the caching component **148** may report (sub-block **409**) the discrepancy to the remote system **104** as a metric. In some embodiments, as shown by sub-block **411**, the caching component **148** may add a confirmation step that is to be performed in real-time if and when the first local NLU result **502** is used for responding to the user speech, such as a confirmation step that asks the user **106** if they want to receive a response despite a low confidence intent resolution (e.g., a low confidence first local NLU result **502**).

(67) If, at block **408**, the caching component **148** determines that the remote NLU result **508** is executable using a local speechlet **146** on the speech interface device **102**, the process **400** may follow the “YES” route from block **408** to block **412**.

(68) At block **412**, the caching component **148** may determine a cacheable status of the remote NLU result **508**. Determining the cacheable status of the remote NLU result **508** can include determining the cacheable status of individual slots of the remote NLU result **508**, which may be based on a cacheable status indicator associated with the individual slots of the remote NLU result **508**. In the example of FIG. 5, the remote NLU result **508** is associated with a cacheable status indicator for at least one slot that indicates the text of the slot (e.g., {Address A}) is cacheable.

(69) At **414** of the process **400**, the caching component **148** may determine whether all slots of the remote NLU result **508** are cacheable (e.g., based on a cacheable status indicator associated with individual slots of the remote NLU result **508**). If, at block **414**, the caching component **148** determines that all slots of the remote NLU result **508** are cacheable, the process **400** may follow the “YES” route from block **414** to block **416**, where the caching component **148** may treat this condition as a correction, and may cache the remote NLU result **508** by storing, in memory of the speech interface device **102**: the remote NLU result **508**, and association data that associates the remote NLU result **508** with the first local ASR result **502**. This is shown in the first cache entry of the cache entries **150** in FIG. 5, where the first local ASR result **502** (e.g., “Take me home now”) is associated with the remote NLU result **508** (e.g., “Navigate to {Address A}”). In some embodiments, an expiration time may be enacted on the stored remote NLU result **508** to expire the remote NLU result **508** from the cache entries **150** (e.g., remove the remote NLU result **508** from local memory) after a period of time since caching the remote NLU result **508**.

(70) If, at block **414**, the caching component **148** determines that the text of at least one slot of the remote NLU result **508** is not cacheable, the process **400** may follow the “NO” route from block **414** to block **418** where the caching component **148** may determine whether any non-cacheable

slots are resolvable locally. If, at block **418**, at least one slot is not resolvable by the local speech processing component **140** performing slot resolution, the process **400** may follow the “NO” route from block **418** to block **420**, where the caching component **148** may treat this condition as a discrepancy. As shown by sub-blocks **419** and **421**, the caching component **148** may report (sub-block **419**) the discrepancy to the remote system **104** as a metric, and may upload (sub-block **421**) one or more local results (e.g., the first local ASR result **502** and/or the first local NLU result **504**) to the remote system **104**. The remote system **104** may use the reported metric and/or the uploaded local results to retrain models (e.g., language models, etc.) used by local speech processing components, such as the local speech processing component **140**, to improve the speech recognition capabilities of the local components.

(71) If, at block **418**, the slot(s) of the remote NLU result **508** is/are resolvable locally using the local speech processing component **140** performing slot resolution on those slots, the process **400** may follow the “YES” route from block **418** to block **422**, where the caching component **148** may treat this condition as a partial correction and may storing the at least one slot that is cacheable, but locally unresolvable, with an unresolved label as part of caching the remote NLU result **508**.

(72) At **424**, following a validation condition, a discrepancy condition, a correction condition, or a partial correction condition, the speech interface device **102** may continue in its response to the user speech by performing an action. This action can be based on the remote directive data **510** received from the remote system **104**. Alternatively, if the hybrid request selector **128** determines to use locally-generated directive data, the action may be based on locally-generated directive data. In the example of FIG. 5, having received a remote response **121** from the remote system **104**, the remote directive data **510** may be used to cause the speech interface device **102** to perform the action, such as outputting content (e.g., audio and/or video content) via an output device **112** of the speech interface device **102**, sending a command to a second device collocated in an environment with the speech interface device **102**, the command instructing an operation to be performed at the second device (e.g., turn on/off a light), and so on.

(73) In some embodiments, the remote directive data **510** (when cacheable) is cached at block **416** or **422** and associated with the first local ASR result **502**. This is shown in the first cache entry of the cache entries **150** in FIG. 5 where the first local ASR result **502** is associated with the remote NLU result **508** and to the remote directive data **510**. Caching the remote directive data **510** may include caching a static response (e.g., a recording of an audible response, such as a speech byte stream, or content that can be played back using the static skill handler **152**). This operation may include determining that a cacheable status indicator (e.g., “C”) associated with the remote directive data **510** indicates that the remote directive data **510** is cacheable.

(74) FIG. 6 is a flow diagram of an example process **600** implemented by a hybrid speech interface device for using cached information for local speech processing. For discussion purposes, the process **600** is described with reference to the previous figures. Furthermore, the process **600** may proceed from either block **230** of the process **200**, or block **424** of the process **400**, as indicated by the off-page reference “B” in FIGS. 2, 4, and 6, where at least some remote speech processing information may be cached and available for use in responding to user speech.

(75) At **602**, the speech interface device may receive second audio data **116** that represents second user speech. In the example of FIGS. 1 and 3, this may be a time in the future when the user **106** utters the expression “Turn off the cellar light” and the remote ASR result **306** (e.g., “Turn off the cellar light”) is already cached in local memory of the speech interface device **102**. In some embodiments, the second audio data **116** may be received at a time when the remote system **104** is unavailable to the speech interface device **102** (e.g., when the Internet is down).

(76) At **604**, the local speech processing component **140** (e.g., the local ASR component **142**) may perform ASR processing on the second audio data **116** to generate a second local ASR result. For example, the second local ASR result may be “Turn off the seller light”, using the running example.

(77) At **606**, the caching component **148** may receive the second local ASR result from the local

speech processing component **140** and may access the association data in the memory of the speech interface device **102** to compare the second local ASR result to cache entries of local ASR results, including, for example, the first local ASR result **302** shown in FIG. 3, and/or the first local ASR result **502** shown in FIG. 5.

(78) At **608**, the caching component **148** may determine that the second local ASR result matches the first local ASR result **302/502**.

(79) At **610**, the caching component **148** may determine whether there is a cached remote NLU result, such as the remote NLU result **508** of FIG. 5, associated with the first local ASR result **502** that matches the second local ASR result. If there is no cached remote NLU result, but there is a cached remote ASR result **306** that is associated with the cached first local ASR result **302**, the process **600** may follow the “NO” route from block **610** to block **612**.

(80) At **612**, the caching component **148** may retrieve the remote ASR result **306** from the memory of the speech interface device **102** as a retrieved remote ASR result **306**.

(81) At **614**, the local speech processing component **140** (e.g., the local NLU component **144**) may perform NLU processing on the retrieved remote ASR result **306** to generate the second local NLU result **312** (as shown in FIG. 3).

(82) At **616**, the speech interface device **102** may perform an action, for a second time (because the first time the speech interface device **102** performed the action, it may not have had a cached remote ASR result to retrieve), based at least in part on the second local NLU result **312**. In some embodiments, the action performed at block **616** may be based on local directive data generated by the local speech processing component **140** based on the second local NLU result **312**.

(83) If, at block **610**, the caching component **148** determines, based on accessing the cache entries **150**, that a remote NLU result **508** is stored in local memory and associated with the first local ASR result **502**, the process **600** may follow the “YES” route from block **610** to block **618**.

(84) At **618**, the caching component **148** may retrieve the remote NLU result **508** from the memory of the speech interface device **102** as a retrieved remote NLU result **508**.

(85) At **620**, if the retrieved remote NLU result **508** was a partial correction that included an unresolved slot(s), the local speech processing component **140** (e.g., the local NLU component **144**) may perform slot resolution to resolve the slot(s) with second text.

(86) At **616**, following block **620** (or following block **618** if there are no unresolved slots in the retrieved remote NLU result **508**), the action may be performed at block **616**, this time, the action being based at least in part on the retrieved remote NLU result **508**, and possibly the retrieved remote NLU result **508** with a slot(s) resolved locally (e.g., with second text for the slot(s) generated by the local speech processing component **140**).

(87) As shown by sub-block **617**, in some embodiments, if remote directive data **510** is cached and is associated with the first local ASR result **502**, the remote directive data **510** may be retrieved (sub-block **617**) from the memory of the speech interface device **102** as retrieved remote directive data **510**, and at block **616**, the action may be performed based on the retrieved remote directive data **510**. In some embodiments, retrieving the remote directive data **510** includes retrieving a static response (e.g., content) from the memory of the speech interface device **102**, and outputting the static response (e.g., content) via an output device **112** of the speech interface device **102** (e.g., using the static skill handler **152**).

(88) As shown by sub-blocks **619** and **621**, in some embodiments, if the speech interface device **102** determines that the retrieved remote NLU result **508** is not executable on the speech interface device, the speech interface device **102** may send (sub-block **619**) the retrieved remote NLU result **508** to the remote system **104**, which may execute on the retrieved remote NLU result **508** to generate second remote directive data, and the speech interface device **102** may receive the second remote directive data from the remote system **102** to perform the action based on the second remote directive data. This allows for “partial outage protection” because if the remote speech processing system **120** is unavailable to the speech interface device **102**, yet the remote system **104** is still

available to execute on an intent, the speech interface device **102** may send the retrieved remote NLU result **508** to the remote system **104**, whose other components (e.g., components other than the remote speech processing system **120**) may be available.

(89) FIG. 7 is a flow diagram of an example process **700** implemented by a hybrid speech interface device **102** for selectively cached remote directive data based on popularity of an utterance or a likelihood of detected utterance in the future. For discussion purposes, the process **700** is described with reference to the previous figures.

(90) At **702**, the speech interface device **102** may receive audio data **116** that represents user speech.

(91) At **704**, the caching component **148** may determine whether the audio data **116** corresponds to an utterance that has been detected more than a threshold number of times, or above a threshold frequency, for a user account that is associated with the speech interface device **102**. In other words, the determination at block **704** is to determine if the utterance is “popular” in that the user **106** has expressed the utterance often in the past. This determination at block **704** may be based on a history of past utterances associated with the user account, which may be maintained in the local memory of the speech interface device **102** (e.g., in the cache entries **150**). If, at block **704**, the caching component **148** determines that the audio data **116** corresponds to an utterance that has been detected more than a threshold number of times, or above a threshold frequency, for a user account that is associated with the speech interface device **102**, the process **700** may follow the “YES” route from block **704** to block **706**, where the caching component **148** may cache remote directive data corresponding to the utterance. For example, in the context of the process **200**, the remote directive data received at block **208** may be cached if it is cacheable (as indicated by a cacheable status indicator), and if it is determined to correspond to a popular utterance at **704** of the process **700**. If, at block **704**, the caching component **148** determines that the audio data **116** corresponds to an utterance that has not been detected more than a threshold number of times, or above a threshold frequency, for a user account that is associated with the speech interface device **102** (e.g., an unpopular utterance), the process **700** may follow the “NO” route from block **704** to block **708**.

(92) At **708**, the caching component **148** may determine whether a probability that the utterance will be detected in the future for a user account that is associated with the speech interface device **102** meets or exceeds a threshold probability. This determination at block **708** may be based on a history of past utterances associated with the user account, and may utilize machine learning to make a prediction as to whether there is a high likelihood that the utterance will be spoken again in the future. If, at block **708**, the caching component **148** determines that a probability that the utterance will be detected in the future for a user account that is associated with the speech interface device **102** meets or exceeds a threshold probability, the process **700** may follow the “YES” route from block **708** to block **706**, where the caching component **148** may cache remote directive data corresponding to the utterance. If, at block **708**, the caching component **148** determines that a probability that the utterance will be detected in the future for a user account that is associated with the speech interface device **102** is less than a threshold probability, the process **700** may follow the “NO” route from block **708** to block **710** where the caching component **148** may refrain from caching remote directive data corresponding to the utterance.

(93) FIG. 8 illustrates example components of an electronic device, such as the hybrid speech interface device **102** of FIG. 1. The speech interface device **102** may be implemented as a standalone device that is relatively simple in terms of functional capabilities with limited input/output components, memory, and processing capabilities. For instance, the speech interface device **102** does not have a keyboard, keypad, or other form of mechanical input. Nor does it have a display (other than simple lights, for instance) or touch screen to facilitate visual presentation and user touch input. Instead, the speech interface device **102** may be implemented with the ability to receive and output audio, a network interface (wireless or wire-based), power, and

processing/memory capabilities. In certain implementations, a limited set of one or more input components may be employed (e.g., a dedicated button to initiate a configuration, power on/off, etc.). Nonetheless, the primary and potentially only mode of user interaction with the speech interface device **102** is through voice input and audible output.

(94) The speech interface device **102** may also be implemented in other form factors, such as a mobile device (e.g., a smart phone or personal digital assistant). The mobile device may include a touch-sensitive display screen and various buttons for providing input as well as additional functionality such as the ability to send and receive telephone calls. Alternative implementations of the speech interface device **102** may also include configuration as a personal computer. The personal computer may include a keyboard, a mouse, a display screen, and any other hardware or functionality that is typically found on a desktop, notebook, netbook, or other personal computing devices. These devices, however, are merely examples and not intended to be limiting, as the techniques described in this disclosure may be used in essentially any device that has an ability to recognize speech input or other types of natural language input.

(95) In the illustrated implementation, the speech interface device **102** includes one or more processors **802** and computer-readable media **804**. In some implementations, the processor(s) **802** may include a central processing unit (CPU), a graphics processing unit (GPU), both CPU and GPU, a microprocessor, a digital signal processor or other processing units or components known in the art. Alternatively, or in addition, the functionally described herein can be performed, at least in part, by one or more hardware logic components. For example, and without limitation, illustrative types of hardware logic components that can be used include field-programmable gate arrays (FPGAs), application-specific integrated circuits (ASICs), application-specific standard products (ASSPs), system-on-a-chip systems (SOCs), complex programmable logic devices (CPLDs), etc. Additionally, each of the processor(s) **802** may possess its own local memory, which also may store program modules, program data, and/or one or more operating systems.

(96) The computer-readable media **804** may include volatile and nonvolatile memory, removable and non-removable media implemented in any method or technology for storage of information, such as computer-readable instructions, data structures, program modules, or other data. Such memory includes, but is not limited to, RAM, ROM, EEPROM, flash memory or other memory technology, CD-ROM, digital versatile disks (DVD) or other optical storage, magnetic cassettes, magnetic tape, magnetic disk storage or other magnetic storage devices, RAID storage systems, or any other medium which can be used to store the desired information and which can be accessed by a computing device. The computer-readable media **804** may be implemented as computer-readable storage media (“CRSM”), which may be any available physical media accessible by the processor(s) **802** to execute instructions stored on the memory **804**. In one basic implementation, CRSM may include random access memory (“RAM”) and Flash memory. In other implementations, CRSM may include, but is not limited to, read-only memory (“ROM”), electrically erasable programmable read-only memory (“EEPROM”), or any other tangible medium which can be used to store the desired information and which can be accessed by the processor(s) **802**.

(97) Several modules such as instruction, datastores, and so forth may be stored within the computer-readable media **804** and configured to execute on the processor(s) **802**. A few example functional modules are shown as applications stored in the computer-readable media **804** and executed on the processor(s) **802**, although the same functionality may alternatively be implemented in hardware, firmware, or as a system on a chip (SOC). At least some of the components shown in FIG. **1** may also be stored in the computer-readable media **804** and executable by the processor(s) **802** to implement the functionality described herein. For example, the voice services component **124**, the hybrid request selector **128** and the local speech processing component **140**, the caching component **148**, as well as their subcomponents, may be stored in the computer-readable media **804** and executable by the processor(s) **802** to implement the

functionality described herein.

(98) An operating system module **806** may be configured to manage hardware within and coupled to the speech interface device **102** for the benefit of other modules. In addition, in some instances the speech interface device **102** may include some or all of one or more secondary-device drivers **808**. In other instances, meanwhile, the speech interface device **102** may be free from the drivers **808** for interacting with second devices. The speech interface device **102** may further include, in some instances, a speech-recognition module **840**, which may correspond to the local speech processing component **140** described with reference to FIG. 1.

(99) The speech interface device **102** may also include a plurality of applications **811** stored in the computer-readable media **804** or otherwise accessible to the speech interface device **102**. In this implementation, the applications **811** are a music player **814**, a movie player **816**, a timer **818**, and a personal shopper **820**. However, the speech interface device **102** may include any number or type of applications and is not limited to the specific examples shown here. The music player **814** may be configured to play songs or other audio files. The movie player **816** may be configured to play movies or other audio visual media. The timer **818** may be configured to provide the functions of a simple timing device and clock. The personal shopper **820** may be configured to assist a user in purchasing items from web-based merchants. When implementing the “hybrid” functionality described herein where a remote system **104** is unavailable to the speech interface device **102**, these applications **811** may be configured to access local resources (e.g., local music or movie libraries, a local shopping list, a local calendar, etc.). In some cases, changes made to these local resources may be synched with remote versions of those resources when the remote system **104** subsequently becomes available to the speech interface device **102**.

(100) Generally, the speech interface device **102** has input devices **822** and output devices **812**. The input devices **822** may include, without limitation, a keyboard, keypad, mouse, touch screen, joystick, control buttons, etc. In some implementations, one or more microphones **810**, which may be similar to the microphone(s) **110** of FIG. 1, may function as input devices **822** to receive audio input, such as user voice input. The output device(s) **812** may be similar to the output device(s) **112** of FIG. 1, and may include, without limitation, a display(s), a light element (e.g., LED), a vibrator to create haptic sensations, or the like. In some implementations, one or more speakers **828** may function as output devices **812** to output audio sounds (e.g., audio content, TTS responses, tones at various frequencies, etc.).

(101) A user **106** may interact with the speech interface device **102** by speaking to it, and the one or more microphone(s) **810** captures the user's speech. The speech interface device **102** can communicate back to the user **106** by emitting audible statements through the speaker **828**. In this manner, the user **106** can interact with the speech interface device **102** solely through speech, without use of a keyboard or display.

(102) The speech interface device **102** may further include a wireless unit **830** coupled to an antenna **832** to facilitate a wireless connection to a network. The wireless unit **830** may implement one or more of various wireless and/or IoT technologies, such as Bluetooth® protocol, Bluetooth Low Energy (BLE) protocol, ZigBee® protocol, Z-wave® protocol, WiFi protocol, and/or any other type of protocol usable to communicate wirelessly between electronic devices in an environment, including those that do and/or do not rely data transmission over the wide area network **118**. As such, the speech interface device **102** may be configured to act as a hub that can communicate with second devices in the environment and control the second devices, such as by using protocol stacks, drivers, and adapters to communicate over a suitable communications protocol. A USB port **834** may further be provided as part of the speech interface device **102** to facilitate a wired connection to a network, or a plug-in network device that communicates with other wireless networks. In addition to the USB port **834**, or as an alternative thereto, other forms of wired connections may be employed, such as a broadband connection, Transmission Control Protocol/Internet Protocol (TCP/IP) protocol connection, etc. The communications interface **114** of

FIG. 1 may include some or all of these components, and/or other components to facilitate communication with other devices.

(103) Accordingly, when implemented as the primarily-voice-operated speech interface device **102**, there may be no input devices, such as navigation buttons, keypads, joysticks, keyboards, touch screens, and the like other than the microphone(s) **810**. Further, there may be no output such as a display for text or graphical output. The speaker(s) **828** may be the main output device. In one implementation, the speech interface device **102** may include non-input control mechanisms, such as basic volume control button(s) for increasing/decreasing volume, as well as power and reset buttons. There may also be a simple light element (e.g., LED) to indicate a state such as, for example, when power is on.

(104) Accordingly, the speech interface device **102** may be implemented as an aesthetically appealing device with smooth and rounded surfaces, with one or more apertures for passage of sound waves. The speech interface device **102** may merely have a power cord and optionally a wired interface (e.g., broadband, USB, etc.). As a result, the speech interface device **102** may be generally produced at a low cost. Once plugged in, the speech interface device **102** may automatically self-configure, or with slight aid of the user, and be ready to use. In other implementations, other I/O components may be added to this basic model, such as specialty buttons, a keypad, display, and the like.

(105) Although the subject matter has been described in language specific to structural features, it is to be understood that the subject matter defined in the appended claims is not necessarily limited to the specific features described. Rather, the specific features are disclosed as illustrative forms of implementing the claims.

Claims

1. A method comprising: generating, by a device, audio data that represents user speech; performing, by the device, speech processing using the audio data to generate a first result; determining, by the device that a remote system is unavailable for processing the first result; determining, by the device and while the remote system is unavailable, that the first result matches with a second result previously stored in a memory of the first device; determining, by the device, that command information associated with the second result is stored in the memory, the command information being previously received from the remote system and being based on previous interactions between the device and the remote system; and performing, by the device, an action specified in the command information, the action being associated with the second result.
2. The method of claim 1, further comprising: generating, by the device, second audio data that represents second user speech; performing the speech processing using the second audio data to determine the action; retrieving the second result from the memory of the device, and performing, by the device and using the second result, the action.
3. The method of claim 1, wherein performing the speech processing comprises performing automatic speech recognition (ASR) processing on the audio data to generate text data.
4. The method of claim 1, wherein the second result includes a model trained by one or more remote components, and wherein performing the speech processing comprises performing the speech processing on the audio data using the model.
5. The method of claim 1, wherein the second result is used to identify a second device collocated in an environment with the device, and further comprising: sending a command to the second device, the command instructing an action to be performed at the second device.
6. The method of claim 1, wherein each of the previous interactions comprises an interaction where a previous user speech is spoken in an environment of the device and a microphone of the device generates respective audio data representing the previous user speech.
7. The method of claim 1, wherein the second result is associated with a previous user speech that

was detected by a second device.

8. The method of claim 7, wherein the second device is collocated in an environment with the device.

9. A device comprising: one or more processors; and memory storing computer-executable instructions that, when executed by the one or more processors, cause the device to: generate audio data that represents user speech; perform speech processing using the audio data to generate a first result; determine that a remote system is unavailable for processing the first result; determine that the first result matches a second result stored in a memory of the device; determine that command information associated with the second result is stored in the memory, the command information being previously received from the remote system and being based on previous interactions between the device and the remote system; and perform an action specified in the command information, the action being associated with the second result.

10. The device of claim 9, wherein the instructions, when executed by the one or more processors, further cause the device to: generate second audio data that represents second user speech; perform the speech processing using the second audio data to determine the action; retrieve the second result from the memory of the device, and perform, using the second result, the action.

11. The device of claim 9, wherein performing the speech processing comprises performing automatic speech recognition (ASR) processing on the audio data to generate text data.

12. The device of claim 9, wherein the second result includes a model trained by one or more remote components, and wherein performing the speech processing comprises performing the speech processing on the audio data using the model.

13. The device of claim 9, wherein the second result is used to identify a second device collocated in an environment with the device, and wherein the instructions, when executed by the one or more processors, further cause the device to: send a command to the second device, the command instructing an action to be performed at the second device.

14. The device of claim 9, wherein each of the previous interactions comprises an interaction where a previous user speech is spoken in an environment of the device and a microphone of the device generates respective audio data representing the previous user speech.

15. The device of claim 9, further comprising a microphone, and wherein each of the previous interactions comprises an interaction where a user spoke in an environment of the device and the microphone generated respective audio data representing the user speech.

16. The device of claim 9, wherein the second result is associated with a previous user speech that was detected by a second device.

17. The device of claim 16, wherein the second device is collocated in an environment with the device.

18. The device of claim 9, further comprising a microphone, and wherein the microphone generated the audio data that represents the user speech.
