

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent

12403919

Kind Code

B2

Date of Patent

September 02, 2025

Inventor(s)

Tsai; Timothy et al.

Hardware fault detection for feedback control systems in autonomous machine applications

Abstract

Systems and methods for detecting hardware faults in computer-based feedback control systems. Multiple instances of the system control program(s) are run on system processors. System sensor data are input to each instance, and the control commands output by each instance are compared. As instantiations of the same programs receive largely the same sensor data, differences between output commands may indicate the presence of one or more hardware faults.

Inventors: Tsai; Timothy (Santa Clara, CA), Jha; Saurabh (Urbana, IL), Hari; Siva (Sunnyvale, CA), Sullivan; Michael (Austin, TX)

Applicant: NVIDIA Corporation (Santa Clara, CA)

Family ID: 80223871

Assignee: NVIDIA Corporation (Santa Clara, CA)

Appl. No.: 18/753841

Filed: June 25, 2024

Prior Publication Data

Document Identifier

Publication Date

US 20240383489 A1

Nov. 21, 2024

Related U.S. Application Data

continuation parent-doc US 16994382 20200814 US 12054164 child-doc US 18753841

Publication Classification

Int. Cl.: B60W50/02 (20120101); B60W60/00 (20200101); G05B23/02 (20060101); G06F11/22 (20060101)

U.S. Cl.:

CPC B60W50/0205 (20130101); B60W60/001 (20200201); G05B23/0262 (20130101); G06F11/22 (20130101);

Field of Classification Search

CPC: B60W (50/0205); B60W (50/0098); B60W (60/001); B60W (60/00); B60W (2050/0006); B60W (2050/021); G05B (23/0262); G05B (19/0428); G06F (11/22)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
10204461	12/2018	Chen	N/A	B60W 50/0205
10885698	12/2020	Muthler et al.	N/A	N/A
11148677	12/2020	Mukai et al.	N/A	N/A
11919529	12/2023	Venkatraman	N/A	B60W 50/085
2004/0019461	12/2003	Bouse et al.	N/A	N/A
2013/0107027	12/2012	Muellhaeuser	348/77	B60K 35/10
2018/0336107	12/2017	Piel et al.	N/A	N/A
2019/0020973	12/2018	Harish et al.	N/A	N/A
2019/0193747	12/2018	Golov	N/A	B60W 50/023
2020/0148201	12/2019	King et al.	N/A	N/A
2020/0379454	12/2019	Trinh et al.	N/A	N/A
2021/0001880	12/2020	Hu	N/A	N/A
2022/0036157	12/2021	Kale et al.	N/A	N/A
2022/0155783	12/2021	Singh	N/A	G05D 1/0088
2022/0161815	12/2021	Van Beek et al.	N/A	N/A
2023/0128700	12/2022	Maes	705/301	G06Q 10/06
2023/0164615	12/2022	Chan et al.	N/A	N/A
2024/0005637	12/2023	Iqbal	N/A	G01C 21/3807

OTHER PUBLICATIONS

IEC 61508, “Functional Safety of Electrical/Electronic/Programmable Electronic Safety-related Systems,” https://en.wikipedia.org/wiki/IEC_61508, accessed on Apr. 1, 2022, 7 pgs. cited by applicant

ISO 26262, “Road vehicle—Functional safety,” International standard for functional safety of electronic system, https://en.wikipedia.org/wiki/ISO_26262, accessed on Sep. 13, 2021, 8 pgs. cited by applicant

“Taxonomy and Definitions for Terms Related to Driving Automation Systems for On-Road Motor Vehicles”, Society of Automotive Engineers (SAE), Standard No. J3016-201609, p. 30 (Sep. 30, 2016). cited by applicant

“Taxonomy and Definitions for Terms Related to Driving Automation Systems for On-Road Motor Vehicles”, Society of Automotive Engineers (SAE), Standard No. J3016-201806, p. 35 (Jun. 15, 2018). cited by applicant

Tsai, et al.; Non-Final Office Action for U.S. Appl. No. 16/994,382, filed Aug. 14, 2020, mailed Jun. 26, 2023, 31 pgs. cited by applicant
Tsai, Tim; Final Office Action for U.S. Appl. No. 16/994,382, filed Aug. 14, 2020, mailed Oct. 10, 2023, 27 pgs. cited by applicant
Tsai, Tlm; Advisory Action for U.S. Appl. No. 16/994,382, filed Aug. 14, 2020, mailed Feb. 21, 2024, 3 pgs. cited by applicant
Tsai, Tim; Notice of Allowance for U.S. Appl. No. 16/994,382, filed Aug. 14, 2020, mailed Mar. 29, 2024, 13 pgs. cited by applicant

Primary Examiner: Kim; Matthew M

Assistant Examiner: Putaraksa; Matthew N

Attorney, Agent or Firm: Taylor Duma L.L.P.

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a continuation of U.S. patent application Ser. No. 16/994,382, filed Aug. 14, 2020, which is herein incorporated by reference in its entirety.

BACKGROUND

(1) Embodiments of the disclosure relate generally to feedback control systems. More specifically, embodiments of the disclosure relate to hardware fault detection to perform feedback control operations in autonomous machine applications.

SUMMARY

(2) Hardware faults occur in, and present challenges for, almost any computer-based system, as they may compromise performance and result in malfunctioning and possibly even unsafe computer-based systems. Of particular difficulty are random hardware faults, which may result from a number of difficult-to-predict phenomena such as neutron or other atomic particle bombardment from the environment, electro-magnetic interference, unexpected voltage drops, heat, and the like. Some of these phenomena, like neutron bombardment, are currently impractical if not impossible to avoid.

(3) While various methods of reducing or compensating for resulting faults exist, they each entail significant drawbacks. For example, hardware design modifications to detect and compensate for faults, such as error correcting codes (ECCs), parity, retry, and the like employ redundant hardware blocks, repeated processing steps, etc., which each incur area, resource, and power overhead and often do not protect every vulnerable hardware element. Even software-based error detection and correction techniques are often only applicable to certain classes of software, and typically require significant manpower to implement, and computing overhead in their execution.

(4) Accordingly, to alleviate challenges associated with hardware faults, systems and methods are described herein for detecting and compensating for hardware faults in computer-based feedback control systems, such as (without limitation), feedback control systems suitable for use with or deployment in autonomous machines and autonomous machine applications. When such control systems employ one or more programs or instruction sets for carrying out feedback control, multiple instances of the same programs are executed to run simultaneously. Sensor output used to determine feedback control commands are input to both instances, and the outputs of the instances are compared. Identical or sufficiently similar output from the multiple instances indicates that the underlying hardware running the instances is sufficiently free of faults, while outputs that differ

significantly may indicate the presence of one or more hardware faults that compromise the execution of at least one instance.

(5) As an example, a feedback control system may direct one or more actuators according to data from one or more sensors that detect various properties of the system and/or its surroundings. This sensor data may thus be fed back to each of multiple instances of a control program for the system, where each instance is executed by system hardware, or circuitry such as parallel processing circuitry. Each instance accordingly generates control commands for the actuators as output resulting from the input sensor data. The presence of hardware faults may be determined by comparing the control commands output by each instance. For example, if the output commands of differing instances differ from each other by greater than some predetermined amount, it may be determined that a hardware fault exists to cause differential operation of the various instances. As another example, one or more machine learning models may be trained to detect faults, such as by training the models to recognize patterns in the control command output of various instances, or patterns in output differences. As another example, multiple outputs of each instance may be considered in the detection of hardware faults, whether by machine learning or other methods. For instance, outputs may be examined over time to determine the presence of persistent errors or patterns in output differences that may indicate certain hardware faults.

(6) Embodiments of the disclosure contemplate simultaneous execution of any number of instances. In some embodiments of the disclosure, processors of a control system may execute two separate instances of their control program, with sensor data being transmitted to each as input, and the output of each compared to determine whether a hardware fault exists. In other embodiments of the disclosure, processors of a control system may execute three or more separate instances of their control program. In such cases, improper outputs, the result of hardware faults, may be determined by majority vote of the various instances. That is, the majority output may be deemed to be the correct output, and the presence of a minority output may be deemed to indicate a hardware fault.

(7) Sensor data may be transmitted to the various system control program instances in any manner. In order to maintain total processor overhead at the same level as that of a single instance, dual instances may be fed sensor data in round robin fashion, in which successive sensor output is fed to alternating instances, e.g., sensor output is switched between instances. In this manner, each instance processes half of the sensor data stream, resulting in each being 50% utilized and the total overhead of both instances being approximately 100% that of a single instance processing all sensor data, i.e., no more overhead than a conventional system. This round robin approach may also be implemented for more than two instances, with sensor data being fed to each different instance in rotating manner. In other embodiments of the disclosure, dual instances may both be fed the same sensor data at half the frequency, i.e., alternating sensor outputs are applied to each instance, with remaining sensor outputs discarded. That is, in one time period both instances are fed the same sensor data, while in the next time period the sensor data is discarded and neither instance is fed any data, and so on. This may also result in approximately 50% utilization for each instance over time, with the total overhead of both instances being approximately 100% that of a single instance processing all sensor data. Other embodiments of the disclosure contemplate any other arrangement of sensor data applied to any number of instances. For example, for improved accuracy, the full sensor data output stream may be fed to both instances, or every instance, so that every instance is fully utilized. This may, however, result in increased processor overhead as compared to a single instance. As another example, one instance may receive the full sensor data stream, while another instance receives less than the full sensor output stream (e.g., half the sensor output, or 50% utilization). Such arrangements may produce more accurate results, albeit via greater processor overhead.

(8) The various instances may be any instances of one or more control programs. Further, instances may be any form of execution of such programs. For example, each instance may be executed by a different virtual device, on one or more processors.

(9) Embodiments of the disclosure contemplate any type of feedback control system. That is, methods and processes of embodiments of the disclosure may be utilized to detect hardware faults in any type or configuration of feedback control system. As one example, embodiments of the disclosure may be employed to detect hardware fault occurrence in processors of an autonomous vehicle or machine.

Description

BRIEF DESCRIPTION OF THE FIGURES

- (1) The above and other objects and advantages of the disclosure will be apparent upon consideration of the following detailed description, taken in conjunction with the accompanying drawings, in which like reference characters refer to like parts throughout, and in which:
- (2) FIG. 1 is a block diagram illustrating operation of an exemplary hardware fault detection system, constructed according to some embodiments of the disclosure;
- (3) FIG. 2 conceptually illustrates various methods of distributing sensor data to the control stacks of FIG. 1, according to some embodiments of the disclosure;
- (4) FIG. 3 is a block diagram of an illustrative feedback control system constructed for use according to some embodiments of the disclosure;
- (5) FIG. 4A is an illustration of an example autonomous vehicle, in accordance with some embodiments of the present disclosure;
- (6) FIG. 4B is an example of camera locations and fields of view for the example autonomous vehicle of FIG. 4A, in accordance with some embodiments of the present disclosure;
- (7) FIG. 4C is a block diagram of an example system architecture for the example autonomous vehicle of FIG. 4A, in accordance with some embodiments of the present disclosure;
- (8) FIG. 4D is a system diagram for communication between cloud-based server(s) and the example autonomous vehicle of FIG. 4A, in accordance with some embodiments of the present disclosure;
- (9) FIG. 5 is a block diagram of an example computing device suitable for use in implementing some embodiments of the present disclosure
- (10) FIGS. 6A-6D each illustrate an exemplary computer system for use in implementing some embodiments of the present disclosure;
- (11) FIGS. 6E-6F illustrate a shared programming model, in accordance with some embodiments of the present disclosure;
- (12) FIGS. 7A-7B each illustrate inference and/or training logic in accordance with some embodiments of the disclosure; and
- (13) FIG. 8 is a flowchart illustrating an exemplary hardware fault detection process, according to some embodiments of the disclosure.

DETAILED DESCRIPTION

- (14) In one embodiment, the disclosure relates to systems and methods for detecting hardware faults in computer-based feedback control systems. Multiple instances of the system control program(s) are run on system processors. System sensor data are input to each instance, and the control commands output by each instance are compared. As instantiations of the same programs receive largely the same sensor data, differences between output commands may indicate the presence of one or more hardware faults.
- (15) The feedback control systems in question may be any feedback control systems. For example, the system may be an autonomous vehicle (AV) whose AV control stack receives feedback from onboard sensors and determines/issues self-driving commands in response. As another example, the system may be a robot whose control stack issues movement commands based on sensor feedback describing its environment. In yet another example, the system may be used in a virtual and/or

simulated environment for testing software and/or hardware components of an autonomous vehicle or robot. Embodiments of the disclosure contemplate any control system capable of implementing multiple instances of its control stack, to govern one or more functions according to feedback from one or more sensors (real or virtual).

(16) FIG. 1 is a block diagram illustrating operation of an exemplary hardware fault detection system, constructed according to embodiments of the disclosure. Here, a number of sensors **100** determine properties of their environment and transmit corresponding data as output to sensor distributor **110**, which routes the sensor output to two or more instances **120**, **130** of a control stack which is configured to control a system. Each instance **120**, **130** receives as input sensor data from one or more sensors **100** via sensor distributor **110**, and generates as output one or more actuation commands directing actuation of the system in response to its detected environment. The output commands generated by the control stack instances **120**, **130** are fed to a checker **140**, which examines the outputs of the instances **120**, **130** to determine whether a hardware fault is present in the processor(s) that execute instances **120**, **130**. If it can be determined which output command corresponds to the hardware fault(s), the checker **140** may, for example, discard the commands corresponding to the hardware fault and forward the commands calculated by properly functioning hardware. Alternatively, if checker **140** cannot determine which output command corresponds to the hardware fault, other actions may be taken, such as raising an alert or alarm, requesting execution of the control stack on a backup system, or the like. If no hardware fault is detected, checker **140** forwards the commands to system actuators **150** for actuation of the feedback control system, as described further below.

(17) The feedback control system may be any system that uses one or more control programs to control one or more actuators in response to input from one or more of its sensors. These one or more control programs, which collectively may be a control stack, can be executed by a processor of the feedback control system. Accordingly, the processor may execute more than one instance of the control stack simultaneously. That is, the processor may execute multiple versions of the same control program so that both versions are run at the same time, to generate actuation commands independently and substantially simultaneously, taking as input the readings or data from the same sensors. As the parallel instances **120**, **130** of the control stack are instances of the same program, and take as input the same or largely the same sensor data, they are expected to generate the same or similar output commands in the absence of any hardware faults. That is, differences in the outputs generated by control stack instances **120** and **130** may be attributed to the presence of hardware faults. Checker **140** may thus determine the presence of hardware faults by whether the commands output by instances **120** and **130** differ by greater than some predetermined amount. If the output commands received by checker **140** differ by an amount deemed sufficiently significant, i.e., if a fault in the system's processors is detected, checker **140** may take various actions to flag the presence of the faults, discard faulty commands, or the like. Checker **140** operation is further described below.

(18) The various instances **120**, **130** may be any instances of one or more control programs. Further, instances may be any form of execution of such programs. For example, any number of virtual devices may be generated on any one or more processors of the feedback control system, with each instance executed by a different virtual machine. As another example, groups of available parallel processor cores may be selected and a different instance may be executed by each group. Any such approach to use of one or more virtual processors may be employed by embodiments of the disclosure. Exemplary suitable virtual processors and their execution are discussed further below.

(19) As above, sensor distributor **110** distributes the output of sensors **100** to each instance **120**, **130** of the system's control stack. Distribution of this data, or data stream, may occur in any manner. FIG. 2 conceptually illustrates various sensor data distribution methods, according to embodiments of the disclosure. In particular, FIG. 2 illustrates four different exemplary methods by which

distributor **110** may distribute sensor data between two different instances **120**, **130**. While each method is illustrated in the context of two control stack instances, it is noted that embodiments of the disclosure contemplate any number of instances, including three or more instances, and the various methods shown may be applied to any number of instances.

(20) Part (a) of FIG. 2 illustrates a first method, which may be referred to as a round robin method, in which data from sensors **100** are transmitted to alternating instances **120** and **130** in successive time periods. Thus, for example, sensor distributor **110** may transmit sensor data output from sensors **100** to only the first instance **120** during time period $t_{sub.0}$, to only the second instance **130** during the next time period $t_{sub.1}$, to only the first instance **120** again during time period $t_{sub.2}$, to only the second instance **130** during time period $t_{sub.3}$, and so on. If sensors **100** output discrete sets of data, such as images, this round robin approach may send successive data sets to alternating instances **120**, **130**. If sensors **100** are sensors that output a continuous data stream, e.g., temperature measurements or the like, the data stream may be directed to differing instances **120**, **130** in each successive time period.

(21) For a single processor executing both instances **120**, **130**, as each instance receives only half the sensor **100** data over time, it is expected that total processor overhead remains substantially the same as if the processor were executing a single instance that received all sensor **100** data. Accordingly, it is expected that this round robin distribution would generate approximately the same processor overhead as a conventional system running a single instance of its control stack. As such, processes of embodiments of the disclosure may be employed to detect hardware faults without significant added processor overhead.

(22) Part (b) of FIG. 2 illustrates a second method, in which data from sensors **100** are transmitted to both instances **120** and **130**, but at only half the normal frequency. More specifically, sensor distributor **110** may transmit sensor data output from sensors **100** to both instances **120** and **130** during time periods $t_{sub.0}$ and $t_{sub.2}$, but would transmit no sensor **100** data during time periods $t_{sub.1}$ and $t_{sub.3}$. Thus, for example, if sensors **100** output discrete sets of data such as images, this method may send every other data set to both instances **120** and **130**, and discard the remaining data sets. If sensors **100** output a continuous data stream, this stream may be directed to both instances **120**, **130** every other time period, and discarded at other times.

(23) It may be observed that method (b) would generate approximately the same processor overhead as a conventional system running a single instance of its control stack, as each instance **120**, **130** receives and processes sensor **100** data half the time. Accordingly, as with method (a), processes employing method (b) may be used to detect hardware faults without added processor overhead as compared to a fully utilized single instance.

(24) Part (c) of FIG. 2 illustrates a third method, in which data from sensors **100** are continuously transmitted to both instances **120** and **130** without interruption. That is, both instances **120** and **130** each receive sensor **100** output in continuous manner without interruption. This approach may be expected to generate more accurate results, as unlike the previous approaches (a) and (b), both instances **120** and **130** receive all sensor **100** data. However, this approach incurs approximately double the memory and processor overhead of a fully utilized single instance, and double the processor overhead of either approach (a) or (b). Nevertheless, this approach may be employed in connection with suitable control systems having sufficient memory and processor resources.

(25) Part (d) of FIG. 2 illustrates a fourth method, in which data from sensors **100** are continuously transmitted to instance **120** without interruption, but only transmitted to instance **130** at half the frequency. That is, while instance **120** receives all sensor **100** data, instance **130** receives sensor **100** data every other time period (or receives every other data set). As instance **120** is fully utilized but instance **130** processes only half the sensor **100** data, this method may incur 50% greater processor overhead than a conventional system running a single instance of its control stack. As instance **120** receives all sensor **100** data, its output may be considered more accurate than the output of instance **130**, which receives half the data from sensor **100**. Accordingly, checker **140**

may transmit the output of instance **120** to system actuators **150**, while the output of instance **130** is used as a check to detect the presence of faults. That is, the system may be run by instance **120**, with instance **130** generating output solely as a check to detect the presence of faults. Accordingly, when a fault is detected, e.g., according to a sufficiently significant difference between the output of instance **130** as compared to that of instance **120**, checker **140** may take appropriate action, such as flagging the presence of a fault, discarding output commands of instance **120**, or the like.

(26) As can be seen from FIG. 2, in certain sensor data distribution methods, instances **120**, **130** do not receive sensor data during every time period. However, actuators **150** may expect instructions every time period. Embodiments of the disclosure contemplate reconciliation of this apparent mismatch in any manner. For example, if instances **120**, **130** do not generate an output command in a particular time period, the previous command may be transmitted again, a command indicating no change from the previous actuator configuration may be transmitted, or the like. Any approach for reconciling such mismatches in actuator command timing may be employed.

(27) FIG. 3 is a block diagram of an illustrative feedback control system constructed for use according to embodiments of the disclosure. Here, the feedback control system comprises a computing device **300**, one or more sensors **310**, and one or more actuators **320**. Computing device **300**, which may be any electronic computing device containing processing circuitry capable of carrying out the hardware fault detection operations of embodiments of the disclosure, is in electronic communication with both a number of sensors **310** and actuators **320**. In operation, sensors **310**, which may correspond to sensors **100** of FIG. 1, may capture and transmit any desired environmental properties to computing device **300**, which then carries out the above described hardware fault detection processes, determining from the output of sensors **310** output commands governing actuators **320**. The computing device **300** transmits these commands to actuators **320**, which actuate the feedback control system.

(28) The feedback control system of FIG. 3 may be any system capable of actuation in response to an input to any of its sensors. As one example, the system of FIG. 3 may be a robotic actuator capable of actuation according to input force or various environmental variables. As another example, the system of FIG. 3 may be an autonomous vehicle capable of self-directed driving in response to a path calculated from its environment and a desired destination, capable of actuation of its propulsion and steering mechanisms in response to driver input, or the like. One such exemplary autonomous vehicle may be autonomous vehicle **400** now described in connection with FIGS. 4A-4D and FIG. 5.

(29) FIG. 4A is an illustration of an example autonomous vehicle **400**, in accordance with some embodiments of the present disclosure. The autonomous vehicle **400** (alternatively referred to herein as the “vehicle **400**”) may include, without limitation, a passenger vehicle, such as a car, a truck, a bus, a first responder vehicle, a shuttle, an electric or motorized bicycle, a motorcycle, a fire truck, a police vehicle, an ambulance, a boat, a construction vehicle, an underwater craft, a drone, and/or another type of vehicle (e.g., that is unmanned and/or that accommodates one or more passengers). Autonomous vehicles are generally described in terms of automation levels, defined by the National Highway Traffic Safety Administration (NHTSA), a division of the US Department of Transportation, and the Society of Automotive Engineers (SAE) “Taxonomy and Definitions for Terms Related to Driving Automation Systems for On-Road Motor Vehicles” (Standard No. J3016-201806, published on Jun. 15, 2018, Standard No. 13016-201609, published on Sep. 30, 2016, and previous and future versions of this standard). The vehicle **400** may be capable of functionality in accordance with one or more of Level 3-Level 5 of the autonomous driving levels. For example, the vehicle **400** may be capable of conditional automation (Level 3), high automation (Level 4), and/or full automation (Level 5), depending on the embodiment.

(30) The vehicle **400** may include components such as a chassis, a vehicle body, wheels (e.g., 2, 4, 6, 8, 18, etc.), tires, axles, and other components of a vehicle. The vehicle **400** may include a propulsion system **450**, such as an internal combustion engine, hybrid electric power plant, an all-

electric engine, and/or another propulsion system type. The propulsion system **450** may be connected to a drive train of the vehicle **400**, which may include a transmission, to enable the propulsion of the vehicle **400**. The propulsion system **450** may be controlled in response to receiving signals from the throttle/accelerator **452**.

(31) A steering system **454**, which may include a steering wheel, may be used to steer the vehicle **400** (e.g., along a desired path or route) when the propulsion system **450** is operating (e.g., when the vehicle is in motion). The steering system **454** may receive signals from a steering actuator **456**. The steering wheel may be optional for full automation (Level 5) functionality.

(32) The brake sensor system **446** may be used to operate the vehicle brakes in response to receiving signals from the brake actuators **448** and/or brake sensors.

(33) Controller(s) **436**, which may include one or more CPU(s), system on chips (SoCs) **404** (FIG. 4C) and/or GPU(s), may provide signals (e.g., representative of commands) to one or more components and/or systems of the vehicle **400**. For example, the controller(s) may send signals to operate the vehicle brakes via one or more brake actuators **448**, to operate the steering system **454** via one or more steering actuators **456**, and/or to operate the propulsion system **450** via one or more throttle/accelerators **452**. The controller(s) **436** may include one or more onboard (e.g., integrated) computing devices (e.g., supercomputers) that process sensor signals, and output operation commands (e.g., signals representing commands) to enable autonomous driving and/or to assist a human driver in driving the vehicle **400**. The controller(s) **436** may include a first controller **436** for autonomous driving functions, a second controller **436** for functional safety functions, a third controller **436** for artificial intelligence functionality (e.g., computer vision), a fourth controller **436** for infotainment functionality, a fifth controller **436** for redundancy in emergency conditions, and/or other controllers. In some examples, a single controller **436** may handle two or more of the above functionalities, two or more controllers **436** may handle a single functionality, and/or any combination thereof.

(34) The controller(s) **436** may provide the signals for controlling one or more components and/or systems of the vehicle **400** in response to sensor data received from one or more sensors (e.g., sensor inputs). The sensor data may be received from, for example and without limitation, global navigation satellite systems sensor(s) **458** (e.g., Global Positioning System sensor(s)), RADAR sensor(s) **460**, ultrasonic sensor(s) **462**, LIDAR sensor(s) **464**, inertial measurement unit (IMU) sensor(s) **466** (e.g., accelerometer(s), gyroscope(s), magnetic compass(es), magnetometer(s), etc.), microphone(s) **496**, stereo camera(s) **468**, wide-view camera(s) **470** (e.g., fisheye cameras), infrared camera(s) **472**, surround camera(s) **474** (e.g., 360 degree cameras), long-range and/or mid-range camera(s) **498**, speed sensor(s) **444** (e.g., for measuring the speed of the vehicle **400**), vibration sensor(s) **442**, steering sensor(s) **440**, brake sensor(s) **446** (e.g., as part of the brake sensor system **446**), and/or other sensor types.

(35) One or more of the controller(s) **436** may receive inputs (e.g., represented by input data) from an instrument cluster **432** of the vehicle **400** and provide outputs (e.g., represented by output data, display data, etc.) via a human-machine interface (HMI) display **434**, an audible annunciator, a loudspeaker, and/or via other components of the vehicle **400**. The outputs may include information such as vehicle velocity, speed, time, map data (e.g., the HD map **422** of FIG. 4C), location data (e.g., the location of the vehicle **400**, such as on a map), direction, location of other vehicles (e.g., an occupancy grid), information about objects and status of objects as perceived by the controller(s) **436**, etc. For example, the HMI display **434** may display information about the presence of one or more objects (e.g., a street sign, caution sign, traffic light changing, etc.), and/or information about driving maneuvers the vehicle has made, is making, or will make (e.g., changing lanes now, taking exit **34B** in two miles, etc.).

(36) The vehicle **400** further includes a network interface **424**, which may use one or more wireless antenna(s) **426** and/or modem(s) to communicate over one or more networks. For example, the network interface **424** may be capable of communication over LTE, WCDMA, UMTS, GSM,

CDMA2000, etc. The wireless antenna(s) **426** may also enable communication between objects in the environment (e.g., vehicles, mobile devices, etc.), using local area network(s), such as Bluetooth, Bluetooth LE, Z-Wave, ZigBee, etc., and/or low power wide-area network(s) (LPWANs), such as LoRaWAN, SigFox, etc.

(37) FIG. **4B** is an example of camera locations and fields of view for the example autonomous vehicle **400** of FIG. **4A**, in accordance with some embodiments of the present disclosure. The cameras and respective fields of view are one example embodiment and are not intended to be limiting. For example, additional and/or alternative cameras may be included and/or the cameras may be located at different locations on the vehicle **400**.

(38) The camera types for the cameras may include, but are not limited to, digital cameras that may be adapted for use with the components and/or systems of the vehicle **400**. The camera(s) may operate at automotive safety integrity level (ASIL) Band/or at another ASIL. The camera types may be capable of any image capture rate, such as 60 frames per second (fps), 120 fps, 240 fps, etc., depending on the embodiment. The cameras may be capable of using rolling shutters, global shutters, another type of shutter, or a combination thereof. In some examples, the color filter array may include a red clear clear clear (RCCC) color filter array, a red clear clear blue (RCCB) color filter array, a red blue green clear (RBGC) color filter array, a Foveon X3 color filter array, a Bayer sensors (RGGB) color filter array, a monochrome sensor color filter array, and/or another type of color filter array. In some embodiments, clear pixel cameras, such as cameras with an RCCC, an RCCB, and/or an RBGC color filter array, may be used in an effort to increase light sensitivity.

(39) In some examples, one or more of the camera(s) may be used to perform advanced driver assistance systems (ADAS) functions (e.g., as part of a redundant or fail-safe design). For example, a Multi-Function Mono Camera may be installed to provide functions including lane departure warning, traffic sign assist and intelligent headlamp control. One or more of the camera(s) (e.g., all of the cameras) may record and provide image data (e.g., video) simultaneously.

(40) One or more of the cameras may be mounted in a mounting assembly, such as a custom-designed (3-D printed) assembly, in order to cut out stray light and reflections from within the car (e.g., reflections from the dashboard reflected in the windshield mirrors) which may interfere with the camera's image data capture abilities. With reference to wing-mirror mounting assemblies, the wing-mirror assemblies may be custom 3-D printed so that the camera mounting plate matches the shape of the wing-mirror. In some examples, the camera(s) may be integrated into the wing-mirror. For side-view cameras, the camera(s) may also be integrated within the four pillars at each corner of the cabin.

(41) Cameras with a field of view that includes portions of the environment in front of the vehicle **400** (e.g., front-facing cameras) may be used for surround view, to help identify forward-facing paths and obstacles, as well aid in, with the help of one or more controllers **436** and/or control SoCs, providing information critical to generating an occupancy grid and/or determining the preferred vehicle paths. Front-facing cameras may be used to perform many of the same ADAS functions as LIDAR, including emergency braking, pedestrian detection, and collision avoidance. Front-facing cameras may also be used for ADAS functions and systems including Lane Departure Warnings (LDW), Autonomous Cruise Control (ACC), and/or other functions such as traffic sign recognition.

(42) A variety of cameras may be used in a front-facing configuration, including, for example, a monocular camera platform that includes a CMOS (complementary metal oxide semiconductor) color imager. Another example may be a wide-view camera(s) **470** that may be used to perceive objects coming into view from the periphery (e.g., pedestrians, crossing traffic or bicycles). Although only one wide-view camera is illustrated in FIG. **4B**, there may any number of wide-view cameras **470** on the vehicle **400**. In addition, long-range camera(s) **498** (e.g., a long-view stereo camera pair) may be used for depth-based object detection, especially for objects for which a neural network has not yet been trained. The long-range camera(s) **498** may also be used for object

detection and classification, as well as basic object tracking.

(43) One or more stereo cameras **468** may also be included in a front-facing configuration. The stereo camera(s) **468** may include an integrated control unit comprising a scalable processing unit, which may provide a programmable logic (e.g., FPGA) and a multi-core micro-processor with an integrated CAN or Ethernet interface on a single chip. Such a unit may be used to generate a 3-D map of the vehicle's environment, including a distance estimate for all the points in the image. An alternative stereo camera(s) **468** may include a compact stereo vision sensor(s) that may include two camera lenses (one each on the left and right) and an image processing chip that may measure the distance from the vehicle to the target object and use the generated information (e.g., metadata) to activate the autonomous emergency braking and lane departure warning functions. Other types of stereo camera(s) **468** may be used in addition to, or alternatively from, those described herein.

(44) Cameras with a field of view that includes portions of the environment to the side of the vehicle **400** (e.g., side-view cameras) may be used for surround view, providing information used to create and update the occupancy grid, as well as to generate side impact collision warnings. For example, surround camera(s) **474** (e.g., four surround cameras **474** as illustrated in FIG. **4B**) may be positioned around the vehicle **400**. The surround camera(s) **474** may include wide-view camera(s) **470**, fisheye camera(s), 360-degree camera(s), and/or the like. For example, four fisheye cameras may be positioned on the vehicle's front, rear, and sides. In an alternative arrangement, the vehicle may use three surround camera(s) **474** (e.g., left, right, and rear), and may leverage one or more other camera(s) (e.g., a forward-facing camera) as a fourth surround-view camera.

(45) Cameras with a field of view that include portions of the environment to the rear of the vehicle **400** (e.g., rear-view cameras) may be used for park assistance, surround view, rear collision warnings, and creating and updating the occupancy grid. A wide variety of cameras may be used including, but not limited to, cameras that are also suitable as a front-facing camera(s) (e.g., long-range and/or mid-range camera(s) **498**, stereo camera(s) **468**), infrared camera(s) **472**, etc.), as described herein.

(46) Cameras with a field of view that include portions of the interior or cabin of vehicle **400** may be used to monitor one or more states of drivers, passengers, or objects in the cabin. Any type of camera may be used including, but not limited to, cabin camera(s) **441**, which may be any type of camera described herein, and which may be placed anywhere on or in vehicle **400** that provides a view of the cabin or interior thereof. For example, cabin camera(s) **441** may be placed within or on some portion of the vehicle **400** dashboard, rear view mirror, side view mirrors, seats, or doors and oriented to capture images of any drivers, passengers, or any other object or portion of the vehicle **400**.

(47) FIG. **4C** is a block diagram of an example system architecture for the example autonomous vehicle **400** of FIG. **4A**, in accordance with some embodiments of the present disclosure. It should be understood that this and other arrangements described herein are set forth only as examples. Other arrangements and elements (e.g., machines, interfaces, functions, orders, groupings of functions, etc.) may be used in addition to or instead of those shown, and some elements may be omitted altogether. Further, many of the elements described herein are functional entities that may be implemented as discrete or distributed components or in conjunction with other components, and in any suitable combination and location. Various functions described herein as being performed by entities may be carried out by hardware, firmware, and/or software. For instance, various functions may be carried out by a processor executing instructions stored in memory.

(48) Each of the components, features, and systems of the vehicle **400** in FIG. **4C** is illustrated as being connected via bus **402**. The bus **402** may include a Controller Area Network (CAN) data interface (alternatively referred to herein as a "CAN bus"). A CAN may be a network inside the vehicle **400** used to aid in control of various features and functionality of the vehicle **400**, such as actuation of brakes, acceleration, braking, steering, windshield wipers, etc. A CAN bus may be configured to have dozens or even hundreds of nodes, each with its own unique identifier (e.g., a

CAN ID). The CAN bus may be read to find steering wheel angle, ground speed, engine revolutions per minute (RPMs), button positions, and/or other vehicle status indicators. The CAN bus may be ASIL B compliant.

(49) Although the bus **402** is described herein as being a CAN bus, this is not intended to be limiting. For example, in addition to, or alternatively from, the CAN bus, FlexRay and/or Ethernet may be used. Additionally, although a single line is used to represent the bus **402**, this is not intended to be limiting. For example, there may be any number of busses **402**, which may include one or more CAN busses, one or more FlexRay busses, one or more Ethernet busses, and/or one or more other types of busses using a different protocol. In some examples, two or more busses **402** may be used to perform different functions, and/or may be used for redundancy. For example, a first bus **402** may be used for collision avoidance functionality and a second bus **402** may be used for actuation control. In any example, each bus **402** may communicate with any of the components of the vehicle **400**, and two or more busses **402** may communicate with the same components. In some examples, each SoC **404**, each controller **436**, and/or each computer within the vehicle may have access to the same input data (e.g., inputs from sensors of the vehicle **400**), and may be connected to a common bus, such the CAN bus.

(50) The vehicle **400** may include one or more controller(s) **436**, such as those described herein with respect to FIG. 4A. The controller(s) **436** may be used for a variety of functions. The controller(s) **436** may be coupled to any of the various other components and systems of the vehicle **400** and may be used for control of the vehicle **400**, artificial intelligence of the vehicle **400**, infotainment for the vehicle **400**, and/or the like.

(51) The vehicle **400** may include a system(s) on a chip (SoC) **404**. The SoC **404** may include CPU(s) **406**, GPU(s) **408**, processor(s) **410**, cache(s) **412**, accelerator(s) **414**, data store(s) **416**, and/or other components and features not illustrated. The SoC(s) **404** may be used to control the vehicle **400** in a variety of platforms and systems. For example, the SoC(s) **404** may be combined in a system (e.g., the system of the vehicle **400**) with an HD map **422** which may obtain map refreshes and/or updates via a network interface **424** from one or more servers (e.g., server(s) **478** of FIG. 4D).

(52) The CPU(s) **406** may include a CPU cluster or CPU complex (alternatively referred to herein as a “CCPLEX”). The CPU(s) **406** may include multiple cores and/or L2 caches. For example, in some embodiments, the CPU(s) **406** may include eight cores in a coherent multi-processor configuration. In some embodiments, the CPU(s) **406** may include four dual-core clusters where each cluster has a dedicated L2 cache (e.g., a 2 MB L2 cache). The CPU(s) **406** (e.g., the CCPLEX) may be configured to support simultaneous cluster operation enabling any combination of the clusters of the CPU(s) **406** to be active at any given time.

(53) The CPU(s) **406** may implement power management capabilities that include one or more of the following features: individual hardware blocks may be clock-gated automatically when idle to save dynamic power; each core clock may be gated when the core is not actively executing instructions due to execution of WFI/WFE instructions; each core may be independently power-gated; each core cluster may be independently clock-gated when all cores are clock-gated or power-gated; and/or each core cluster may be independently power-gated when all cores are power-gated. The CPU(s) **406** may further implement an enhanced algorithm for managing power states, where allowed power states and expected wakeup times are specified, and the hardware/microcode determines the best power state to enter for the core, cluster, and CCPLEX. The processing cores may support simplified power state entry sequences in software with the work offloaded to microcode.

(54) The GPU(s) **408** may include an integrated GPU (alternatively referred to herein as an “iGPU”). The GPU(s) **408** may be programmable and may be efficient for parallel workloads. The GPU(s) **408**, in some examples, may use an enhanced tensor instruction set. The GPU(s) **408** may include one or more streaming microprocessors, where each streaming microprocessor may include

an L1 cache (e.g., an L1 cache with at least 96 KB storage capacity), and two or more of the streaming microprocessors may share an L2 cache (e.g., an L2 cache with a 512 KB storage capacity). In some embodiments, the GPU(s) **408** may include at least eight streaming microprocessors. The GPU(s) **408** may use computer-based application programming interface(s) (API(s)). In addition, the GPU(s) **408** may use one or more parallel computing platforms and/or programming models (e.g., NVIDIA's CUDA).

(55) The GPU(s) **408** may be power-optimized for best performance in automotive and embedded use cases. For example, the GPU(s) **408** may be fabricated on a Fin field-effect transistor (FinFET). However, this is not intended to be limiting, and the GPU(s) **408** may be fabricated using other semiconductor manufacturing processes. Each streaming microprocessor may incorporate a number of mixed-precision processing cores partitioned into multiple blocks. For example, and without limitation, 64 PF32 cores and 32 PF64 cores may be partitioned into four processing blocks. In such an example, each processing block may be allocated 16 FP32 cores, 8 FP64 cores, 16 INT32 cores, two mixed-precision NVIDIA TENSOR COREs for deep learning matrix arithmetic, an L0 instruction cache, a warp scheduler, a dispatch unit, and/or a 64 KB register file. In addition, the streaming microprocessors may include independent parallel integer and floating-point data paths to provide for efficient execution of workloads with a mix of computation and addressing calculations. The streaming microprocessors may include independent thread-scheduling capability to enable finer-grain synchronization and cooperation between parallel threads. The streaming microprocessors may include a combined L1 data cache and shared memory unit in order to improve performance while simplifying programming.

(56) The GPU(s) **408** may include a high bandwidth memory (HBM) and/or a 16 GB HBM2 memory subsystem to provide, in some examples, about 900 GB/second peak memory bandwidth. In some examples, in addition to, or alternatively from, the HBM memory, a synchronous graphics random-access memory (SGRAM) may be used, such as a graphics double data rate type five synchronous random-access memory (GDDR5).

(57) The GPU(s) **408** may include unified memory technology including access counters to allow for more accurate migration of memory pages to the processor that accesses them most frequently, thereby improving efficiency for memory ranges shared between processors. In some examples, address translation services (ATS) support may be used to allow the GPU(s) **408** to access the CPU(s) **406** page tables directly. In such examples, when the GPU(s) **408** memory management unit (MMU) experiences a miss, an address translation request may be transmitted to the CPU(s) **406**. In response, the CPU(s) **406** may look in its page tables for the virtual-to-physical mapping for the address and transmits the translation back to the GPU(s) **408**. As such, unified memory technology may allow a single unified virtual address space for memory of both the CPU(s) **406** and the GPU(s) **408**, thereby simplifying the GPU(s) **408** programming and porting of applications to the GPU(s) **408**.

(58) In addition, the GPU(s) **408** may include an access counter that may keep track of the frequency of access of the GPU(s) **408** to memory of other processors. The access counter may help ensure that memory pages are moved to the physical memory of the processor that is accessing the pages most frequently.

(59) The SoC(s) **404** may include any number of cache(s) **412**, including those described herein. For example, the cache(s) **412** may include an L3 cache that is available to both the CPU(s) **406** and the GPU(s) **408** (e.g., that is connected to both the CPU(s) **406** and the GPU(s) **408**). The cache(s) **412** may include a write-back cache that may keep track of states of lines, such as by using a cache coherence protocol (e.g., MEI, MESI, MSI, etc.). The L3 cache may include 4 MB or more, depending on the embodiment, although smaller cache sizes may be used.

(60) The SoC(s) **404** may include an arithmetic logic unit(s) (ALU(s)) which may be leveraged in performing processing with respect to any of the variety of tasks or operations of the vehicle **400**—such as processing DNNs. In addition, the SoC(s) **404** may include a floating point unit(s) (FPU(s))

—or other math coprocessor or numeric coprocessor types—for performing mathematical operations within the system. For example, the SoC(s) **104** may include one or more FPUs integrated as execution units within a CPU(s) **406** and/or GPU(s) **408**.

(61) The SoC(s) **404** may include one or more accelerators **414** (e.g., hardware accelerators, software accelerators, or a combination thereof). For example, the SoC(s) **404** may include a hardware acceleration cluster that may include optimized hardware accelerators and/or large on-chip memory. The large on-chip memory (e.g., 4 MB of SRAM), may enable the hardware acceleration cluster to accelerate neural networks and other calculations. The hardware acceleration cluster may be used to complement the GPU(s) **408** and to off-load some of the tasks of the GPU(s) **408** (e.g., to free up more cycles of the GPU(s) **408** for performing other tasks). As an example, the accelerator(s) **414** may be used for targeted workloads (e.g., perception, convolutional neural networks (CNNs), etc.) that are stable enough to be amenable to acceleration. The term “CNN,” as used herein, may include all types of CNNs, including region-based or regional convolutional neural networks (RCNNs) and Fast RCNNs (e.g., as used for object detection).

(62) The accelerator(s) **414** (e.g., the hardware acceleration cluster) may include a deep learning accelerator(s) (DLA). The DLA(s) may include one or more Tensor processing units (TPUs) that may be configured to provide an additional ten trillion operations per second for deep learning applications and inferencing. The TPUs may be accelerators configured to, and optimized for, performing image processing functions (e.g., for CNNs, RCNNs, etc.). The DLA(s) may further be optimized for a specific set of neural network types and floating point operations, as well as inferencing. The design of the DLA(s) may provide more performance per millimeter than a general-purpose GPU, and vastly exceeds the performance of a CPU. The TPU(s) may perform several functions, including a single-instance convolution function, supporting, for example, INT8, INT16, and FP16 data types for both features and weights, as well as post-processor functions.

(63) The DLA(s) may quickly and efficiently execute neural networks, especially CNNs, on processed or unprocessed data for any of a variety of functions, including, for example and without limitation: a CNN for object identification and detection using data from camera sensors; a CNN for distance estimation using data from camera sensors; a CNN for emergency vehicle detection and identification and detection using data from microphones; a CNN for facial recognition and vehicle owner identification using data from camera sensors; and/or a CNN for security and/or safety related events.

(64) The DLA(s) may perform any function of the GPU(s) **408**, and by using an inference accelerator, for example, a designer may target either the DLA(s) or the GPU(s) **408** for any function. For example, the designer may focus processing of CNNs and floating point operations on the DLA(s) and leave other functions to the GPU(s) **408** and/or other accelerator(s) **414**.

(65) The accelerator(s) **414** (e.g., the hardware acceleration cluster) may include a programmable vision accelerator(s) (PVA), which may alternatively be referred to herein as a computer vision accelerator. The PVA(s) may be designed and configured to accelerate computer vision algorithms for the advanced driver assistance systems (ADAS), autonomous driving, and/or augmented reality (AR) and/or virtual reality (VR) applications. The PVA(s) may provide a balance between performance and flexibility. For example, each PVA(s) may include, for example and without limitation, any number of reduced instruction set computer (RISC) cores, direct memory access (DMA), and/or any number of vector processors.

(66) The RISC cores may interact with image sensors (e.g., the image sensors of any of the cameras described herein), image signal processor(s), and/or the like. Each of the RISC cores may include any amount of memory. The RISC cores may use any of a number of protocols, depending on the embodiment. In some examples, the RISC cores may execute a real-time operating system (RTOS). The RISC cores may be implemented using one or more integrated circuit devices, application specific integrated circuits (ASICs), and/or memory devices. For example, the RISC cores may include an instruction cache and/or a tightly coupled RAM.

(67) The DMA may enable components of the PVA(s) to access the system memory independently of the CPU(s) **406**. The DMA may support any number of features used to provide optimization to the PVA including, but not limited to, supporting multi-dimensional addressing and/or circular addressing. In some examples, the DMA may support up to six or more dimensions of addressing, which may include block width, block height, block depth, horizontal block stepping, vertical block stepping, and/or depth stepping.

(68) The vector processors may be programmable processors that may be designed to efficiently and flexibly execute programming for computer vision algorithms and provide signal processing capabilities. In some examples, the PVA may include a PVA core and two vector processing subsystem partitions. The PVA core may include a processor subsystem, DMA engine(s) (e.g., two DMA engines), and/or other peripherals. The vector processing subsystem may operate as the primary processing engine of the PVA, and may include a vector processing unit (VPU), an instruction cache, and/or vector memory (e.g., VMEM). A VPU core may include a digital signal processor such as, for example, a single instruction, multiple data (SIMD), very long instruction word (VLIW) digital signal processor. The combination of the SIMD and VLIW may enhance throughput and speed.

(69) Each of the vector processors may include an instruction cache and may be coupled to dedicated memory. As a result, in some examples, each of the vector processors may be configured to execute independently of the other vector processors. In other examples, the vector processors that are included in a particular PVA may be configured to employ data parallelism. For example, in some embodiments, the plurality of vector processors included in a single PVA may execute the same computer vision algorithm, but on different regions of an image. In other examples, the vector processors included in a particular PVA may simultaneously execute different computer vision algorithms, on the same image, or even execute different algorithms on sequential images or portions of an image. Among other things, any number of PVAs may be included in the hardware acceleration cluster and any number of vector processors may be included in each of the PVAs. In addition, the PVA(s) may include additional error correcting code (ECC) memory, to enhance overall system safety.

(70) The accelerator(s) **414** (e.g., the hardware acceleration cluster) may include a computer vision network on-chip and SRAM, for providing a high-bandwidth, low latency SRAM for the accelerator(s) **414**. In some examples, the on-chip memory may include at least 4 MB SRAM, consisting of, for example and without limitation, eight field-configurable memory blocks, that may be accessible by both the PVA and the DLA. Each pair of memory blocks may include an advanced peripheral bus (APB) interface, configuration circuitry, a controller, and a multiplexer. Any type of memory may be used. The PVA and DLA may access the memory via a backbone that provides the PVA and DLA with high-speed access to memory. The backbone may include a computer vision network on-chip that interconnects the PVA and the DLA to the memory (e.g., using the APB).

(71) The computer vision network on-chip may include an interface that determines, before transmission of any control signal/address/data, that both the PVA and the DLA provide ready and valid signals. Such an interface may provide for separate phases and separate channels for transmitting control signals/addresses/data, as well as burst-type communications for continuous data transfer. This type of interface may comply with ISO 26262 or IEC 61508 standards, although other standards and protocols may be used.

(72) In some examples, the SoC(s) **404** may include a real-time ray-tracing hardware accelerator, such as described in U.S. patent application Ser. No. 16/101,232, filed on Aug. 10, 2018. The real-time ray-tracing hardware accelerator may be used to quickly and efficiently determine the positions and extents of objects (e.g., within a world model), to generate real-time visualization simulations, for RADAR signal interpretation, for sound propagation synthesis and/or analysis, for simulation of SONAR systems, for general wave propagation simulation, for comparison to

LIDAR data for purposes of localization and/or other functions, and/or for other uses. In some embodiments, one or more tree traversal units (TTUs) may be used for executing one or more ray-tracing related operations.

(73) The accelerator(s) **414** (e.g., the hardware accelerator cluster) have a wide array of uses for autonomous driving. The PVA may be a programmable vision accelerator that may be used for key processing stages in ADAS and autonomous vehicles. The PVA's capabilities are a good match for algorithmic domains needing predictable processing, at low power and low latency. In other words, the PVA performs well on semi-dense or dense regular computation, even on small data sets, which need predictable run-times with low latency and low power. Thus, in the context of platforms for autonomous vehicles, the PVAs are designed to run classic computer vision algorithms, as they are efficient at object detection and operating on integer math.

(74) For example, according to one embodiment of the technology, the PVA is used to perform computer stereo vision. A semi-global matching-based algorithm may be used in some examples, although this is not intended to be limiting. Many applications for Level 3-5 autonomous driving require motion estimation/stereo matching on-the-fly (e.g., structure from motion, pedestrian recognition, lane detection, etc.). The PVA may perform computer stereo vision function on inputs from two monocular cameras.

(75) In some examples, the PVA may be used to perform dense optical flow. For example, the PVA may be used to process raw RADAR data (e.g., using a 4D Fast Fourier Transform) to provide a processed RADAR signal before emitting the next RADAR pulse. In other examples, the PVA is used for time of flight depth processing, by processing raw time of flight data to provide processed time of flight data, for example.

(76) The DLA may be used to run any type of network to enhance control and driving safety, including, for example, a neural network that outputs a measure of confidence for each object detection. Such a confidence value may be interpreted as a probability, or as providing a relative “weight” of each detection compared to other detections. This confidence value enables the system to make further decisions regarding which detections should be considered as true positive detections rather than false positive detections. For example, the system may set a threshold value for the confidence and consider only the detections exceeding the threshold value as true positive detections. In an automatic emergency braking (AEB) system, false positive detections would cause the vehicle to automatically perform emergency braking, which is obviously undesirable. Therefore, only the most confident detections should be considered as triggers for AEB. The DLA may run a neural network for regressing the confidence value. The neural network may take as its input at least some subset of parameters, such as bounding box dimensions, ground plane estimate obtained (e.g. from another subsystem), inertial measurement unit (IMU) sensor **466** output that correlates with the vehicle **400** orientation, distance, 3D location estimates of the object obtained from the neural network and/or other sensors (e.g., LIDAR sensor(s) **464** or RADAR sensor(s) **460**), among others.

(77) The SoC(s) **404** may include data store(s) **416** (e.g., memory). The data store(s) **416** may be on-chip memory of the SoC(s) **404**, which may store neural networks to be executed on the GPU and/or the DLA. In some examples, the data store(s) **416** may be large enough in capacity to store multiple instances of neural networks for redundancy and safety. The data store(s) **416** may comprise L2 or L3 cache(s) **412**. Reference to the data store(s) **416** may include reference to the memory associated with the PVA, DLA, and/or other accelerator(s) **414**, as described herein.

(78) The SoC(s) **404** may include one or more processor(s) **410** (e.g., embedded processors). The processor(s) **410** may include a boot and power management processor that may be a dedicated processor and subsystem to handle boot power and management functions and related security enforcement. The boot and power management processor may be a part of the SoC(s) **404** boot sequence and may provide runtime power management services. The boot power and management processor may provide clock and voltage programming, assistance in system low power state

transitions, management of SoC(s) **404** thermal and temperature sensors, and/or management of the SoC(s) **404** power states. Each temperature sensor may be implemented as a ring-oscillator whose output frequency is proportional to temperature, and the SoC(s) **404** may use the ring-oscillators to detect temperatures of the CPU(s) **406**, GPU(s) **408**, and/or accelerator(s) **414**. If temperatures are determined to exceed a threshold, the boot and power management processor may enter a temperature fault routine and put the SoC(s) **404** into a lower power state and/or put the vehicle **400** into a chauffeur to safe-stop mode (e.g., bring the vehicle **400** to a safe stop).

(79) The processor(s) **410** may further include a set of embedded processors that may serve as an audio processing engine. The audio processing engine may be an audio subsystem that enables full hardware support for multi-channel audio over multiple interfaces, and a broad and flexible range of audio I/O interfaces. In some examples, the audio processing engine is a dedicated processor core with a digital signal processor with dedicated RAM.

(80) The processor(s) **410** may further include an always-on processor engine that may provide necessary hardware features to support low power sensor management and wake use cases. The always-on processor engine may include a processor core, a tightly coupled RAM, supporting peripherals (e.g., timers and interrupt controllers), various I/O controller peripherals, and routing logic.

(81) The processor(s) **410** may further include a safety cluster engine that includes a dedicated processor subsystem to handle safety management for automotive applications. The safety cluster engine may include two or more processor cores, a tightly coupled RAM, support peripherals (e.g., timers, an interrupt controller, etc.), and/or routing logic. In a safety mode, the two or more cores may operate in a lockstep mode and function as a single core with comparison logic to detect any differences between their operations.

(82) The processor(s) **410** may further include a real-time camera engine that may include a dedicated processor subsystem for handling real-time camera management.

(83) The processor(s) **410** may further include a high dynamic range signal processor that may include an image signal processor that is a hardware engine that is part of the camera processing pipeline.

(84) The processor(s) **410** may include a video image compositor that may be a processing block (e.g., implemented on a microprocessor) that implements video post-processing functions needed by a video playback application to produce the final image for the player window. The video image compositor may perform lens distortion correction on wide-view camera(s) **470**, surround camera(s) **474**, and/or on in-cabin monitoring camera sensors. An in-cabin monitoring camera sensor is preferably monitored by a neural network running on another instance of the advanced SoC, configured to identify in-cabin events and respond accordingly. An in-cabin system may perform lip reading to activate cellular service and place a phone call, dictate emails, change the vehicle's destination, activate or change the vehicle's infotainment system and settings, or provide voice-activated web surfing. Certain functions are available to the driver only when the vehicle is operating in an autonomous mode, and are disabled otherwise.

(85) The video image compositor may include enhanced temporal noise reduction for both spatial and temporal noise reduction. For example, where motion occurs in a video, the noise reduction weights spatial information appropriately, decreasing the weight of information provided by adjacent frames. Where an image or portion of an image does not include motion, the temporal noise reduction performed by the video image compositor may use information from the previous image to reduce noise in the current image.

(86) The video image compositor may also be configured to perform stereo rectification on input stereo lens frames. The video image compositor may further be used for user interface composition when the operating system desktop is in use, and the GPU(s) **408** is not required to continuously render new surfaces. Even when the GPU(s) **408** is powered on and actively performing 3D rendering, the video image compositor may be used to offload the GPU(s) **408** to improve

performance and responsiveness.

(87) The SoC(s) **404** may further include a mobile industry processor interface (MIPI) camera serial interface for receiving video and input from cameras, a high-speed interface, and/or a video input block that may be used for camera and related pixel input functions. The SoC(s) **404** may further include an input/output controller(s) that may be controlled by software and may be used for receiving I/O signals that are uncommitted to a specific role.

(88) The SoC(s) **404** may further include a broad range of peripheral interfaces to enable communication with peripherals, audio codecs, power management, and/or other devices. The SoC(s) **404** may be used to process data from cameras (e.g., connected over Gigabit Multimedia Serial Link and Ethernet), sensors (e.g., LIDAR sensor(s) **464**, RADAR sensor(s) **460**, etc. that may be connected over Ethernet), data from bus **402** (e.g., speed of vehicle **400**, steering wheel position, etc.), data from GNSS sensor(s) **458** (e.g., connected over Ethernet or CAN bus). The SoC(s) **404** may further include dedicated high-performance mass storage controllers that may include their own DMA engines, and that may be used to free the CPU(s) **406** from routine data management tasks.

(89) The SoC(s) **404** may be an end-to-end platform with a flexible architecture that spans automation levels 3-5, thereby providing a comprehensive functional safety architecture that leverages and makes efficient use of computer vision and ADAS techniques for diversity and redundancy, provides a platform for a flexible, reliable driving software stack, along with deep learning tools. The SoC(s) **404** may be faster, more reliable, and even more energy-efficient and space-efficient than conventional systems. For example, the accelerator(s) **414**, when combined with the CPU(s) **406**, the GPU(s) **408**, and the data store(s) **416**, may provide for a fast, efficient platform for level 3-5 autonomous vehicles.

(90) The technology thus provides capabilities and functionality that cannot be achieved by conventional systems. For example, computer vision algorithms may be executed on CPUs, which may be configured using high-level programming language, such as the C programming language, to execute a wide variety of processing algorithms across a wide variety of visual data. However, CPUs are oftentimes unable to meet the performance requirements of many computer vision applications, such as those related to execution time and power consumption, for example. In particular, many CPUs are unable to execute complex object detection algorithms in real-time, which is a requirement of in-vehicle ADAS applications, and a requirement for practical Level 3-5 autonomous vehicles.

(91) In contrast to conventional systems, by providing a CPU complex, GPU complex, and a hardware acceleration cluster, the technology described herein allows for multiple neural networks to be performed simultaneously and/or sequentially, and for the results to be combined together to enable Level 3-5 autonomous driving functionality. For example, a CNN executing on the DLA or dGPU (e.g., the GPU(s) **420**) may include a text and word recognition, allowing the supercomputer to read and understand traffic signs, including signs for which the neural network has not been specifically trained. The DLA may further include a neural network that is able to identify, interpret, and provide semantic understanding of the sign, and to pass that semantic understanding to the path-planning modules running on the CPU Complex.

(92) As another example, multiple neural networks may be run simultaneously, as is required for Level 3, 4, or 5 driving. For example, a warning sign consisting of “Caution: flashing lights indicate icy conditions,” along with an electric light, may be independently or collectively interpreted by several neural networks. The sign itself may be identified as a traffic sign by a first deployed neural network (e.g., a neural network that has been trained), the text “Flashing lights indicate icy conditions” may be interpreted by a second deployed neural network, which informs the vehicle's path-planning software (preferably executing on the CPU Complex) that when flashing lights are detected, icy conditions exist. The flashing light may be identified by operating a third deployed neural network over multiple frames, informing the vehicle's path-planning software

of the presence (or absence) of flashing lights. All three neural networks may run simultaneously, such as within the DLA and/or on the GPU(s) **408**.

(93) In some examples, a CNN for facial recognition and vehicle owner identification may use data from camera sensors to identify the presence of an authorized driver and/or owner of the vehicle **400**. The always-on sensor processing engine may be used to unlock the vehicle when the owner approaches the driver door and turn on the lights, and, in security mode, to disable the vehicle when the owner leaves the vehicle. In this way, the SoC(s) **404** provide for security against theft and/or carjacking.

(94) In another example, a CNN for emergency vehicle detection and identification may use data from microphones **496** to detect and identify emergency vehicle sirens. In contrast to conventional systems, which use general classifiers to detect sirens and manually extract features, the SoC(s) **404** use the CNN for classifying environmental and urban sounds, as well as classifying visual data. In a preferred embodiment, the CNN running on the DLA is trained to identify the relative closing speed of the emergency vehicle (e.g., by using the Doppler Effect).

(95) The CNN may also be trained to identify emergency vehicles specific to the local area in which the vehicle is operating, as identified by GNSS sensor(s) **458**. Thus, for example, when operating in Europe the CNN will seek to detect European sirens, and when in the United States the CNN will seek to identify only North American sirens. Once an emergency vehicle is detected, a control program may be used to execute an emergency vehicle safety routine, slowing the vehicle, pulling over to the side of the road, parking the vehicle, and/or idling the vehicle, with the assistance of ultrasonic sensors **462**, until the emergency vehicle(s) passes.

(96) The vehicle may include a CPU(s) **418** (e.g., discrete CPU(s), or dCPU(s)), that may be coupled to the SoC(s) **404** via a high-speed interconnect (e.g., PCIe). The CPU(s) **418** may include an X86 processor, for example. The CPU(s) **418** may be used to perform any of a variety of functions, including arbitrating potentially inconsistent results between ADAS sensors and the SoC(s) **404**, and/or monitoring the status and health of the controller(s) **436** and/or infotainment SoC **430**, for example.

(97) The vehicle **400** may include a GPU(s) **420** (e.g., discrete GPU(s), or dGPU(s)), that may be coupled to the SoC(s) **404** via a high-speed interconnect (e.g., NVIDIA's NVLINK). The GPU(s) **420** may provide additional artificial intelligence functionality, such as by executing redundant and/or different neural networks, and may be used to train and/or update neural networks based on input (e.g., sensor data) from sensors of the vehicle **400**.

(98) The vehicle **400** may further include the network interface **424** which may include one or more wireless antennas **426** (e.g., one or more wireless antennas for different communication protocols, such as a cellular antenna, a Bluetooth antenna, etc.). The network interface **424** may be used to enable wireless connectivity over the Internet with the cloud (e.g., with the server(s) **478** and/or other network devices), with other vehicles, and/or with computing devices (e.g., client devices of passengers). To communicate with other vehicles, a direct link may be established between the two vehicles and/or an indirect link may be established (e.g., across networks and over the Internet). Direct links may be provided using a vehicle-to-vehicle communication link. The vehicle-to-vehicle communication link may provide the vehicle **400** information about vehicles in proximity to the vehicle **400** (e.g., vehicles in front of, on the side of, and/or behind the vehicle **400**). This functionality may be part of a cooperative adaptive cruise control functionality of the vehicle **400**.

(99) The network interface **424** may include an SoC that provides modulation and demodulation functionality and enables the controller(s) **436** to communicate over wireless networks. The network interface **424** may include a radio frequency front-end for up-conversion from baseband to radio frequency, and down conversion from radio frequency to baseband. The frequency conversions may be performed through well-known processes, and/or may be performed using super-heterodyne processes. In some examples, the radio frequency front end functionality may be provided by a separate chip. The network interface may include wireless functionality for

communicating over LTE, WCDMA, UMTS, GSM, CDMA2000, Bluetooth, Bluetooth LE, Wi-Fi, Z-Wave, ZigBee, LoRaWAN, and/or other wireless protocols.

(100) The vehicle **400** may further include data store(s) **428**, which may include off-chip (e.g., off the SoC(s) **404**) storage. The data store(s) **428** may include one or more storage elements including RAM, SRAM, DRAM, VRAM, Flash, hard disks, and/or other components and/or devices that may store at least one bit of data.

(101) The vehicle **400** may further include GNSS sensor(s) **458** (e.g., GPS and/or assisted GPS sensors), to assist in mapping, perception, occupancy grid generation, and/or path planning functions. Any number of GNSS sensor(s) **458** may be used, including, for example and without limitation, a GPS using a USB connector with an Ethernet to serial (RS-232) bridge. The vehicle **400** may further include RADAR sensor(s) **460**. The RADAR sensor(s) **460** may be used by the vehicle **400** for long-range vehicle detection, even in darkness and/or severe weather conditions. RADAR functional safety levels may be ASIL B. The RADAR sensor(s) **460** may use the CAN and/or the bus **402** (e.g., to transmit data generated by the RADAR sensor(s) **460**) for control and to access object tracking data, with access to Ethernet to access raw data, in some examples. A wide variety of RADAR sensor types may be used. For example, and without limitation, the RADAR sensor(s) **460** may be suitable for front, rear, and side RADAR use. In some example, Pulse Doppler RADAR sensor(s) are used.

(102) The RADAR sensor(s) **460** may include different configurations, such as long-range with narrow field of view, short-range with wide field of view, short-range side coverage, etc. In some examples, long-range RADAR may be used for adaptive cruise control functionality. The long-range RADAR systems may provide a broad field of view realized by two or more independent scans, such as within a 250 m range. The RADAR sensor(s) **460** may help in distinguishing between static and moving objects, and may be used by ADAS systems for emergency brake assist and forward collision warning. Long-range RADAR sensors may include monostatic multimodal RADAR with multiple (e.g., six or more) fixed RADAR antennae and a high-speed CAN and FlexRay interface. In an example with six antennae, the central four antennae may create a focused beam pattern, designed to record the surrounding of the vehicle **400** at higher speeds with minimal interference from traffic in adjacent lanes. The other two antennae may expand the field of view, making it possible to quickly detect vehicles entering or leaving the vehicle's **400** lane.

(103) Mid-range RADAR systems may include, as an example, a range of up to 460 m (front) or 80 m (rear), and a field of view of up to 42 degrees (front) or 450 degrees (rear). Short-range RADAR systems may include, without limitation, RADAR sensors designed to be installed at both ends of the rear bumper. When installed at both ends of the rear bumper, such a RADAR sensor system may create two beams that constantly monitor the blind spot in the rear and next to the vehicle.

(104) Short-range RADAR systems may be used in an ADAS system for blind spot detection and/or lane change assist.

(105) The vehicle **400** may further include ultrasonic sensor(s) **462**. The ultrasonic sensor(s) **462**, which may be positioned at the front, back, and/or the sides of the vehicle **400**, may be used for park assist and/or to create and update an occupancy grid. A wide variety of ultrasonic sensor(s) **462** may be used, and different ultrasonic sensor(s) **462** may be used for different ranges of detection (e.g., 2.5 m, 4 m). The ultrasonic sensor(s) **462** may operate at functional safety levels of ASIL B.

(106) The vehicle **400** may include LIDAR sensor(s) **464**. The LIDAR sensor(s) **464** may be used for object and pedestrian detection, emergency braking, collision avoidance, and/or other functions. The LIDAR sensor(s) **464** may be functional safety level ASIL B. In some examples, the vehicle **400** may include multiple LIDAR sensors **464** (e.g., two, four, six, etc.) that may use Ethernet (e.g., to provide data to a Gigabit Ethernet switch).

(107) In some examples, the LIDAR sensor(s) **464** may be capable of providing a list of objects and their distances for a 360-degree field of view. Commercially available LIDAR sensor(s) **464**

may have an advertised range of approximately 100 m, with an accuracy of 2 cm-3 cm, and with support for a 100 Mbps Ethernet connection, for example. In some examples, one or more non-protruding LIDAR sensors **464** may be used. In such examples, the LIDAR sensor(s) **464** may be implemented as a small device that may be embedded into the front, rear, sides, and/or corners of the vehicle **400**. The LIDAR sensor(s) **464**, in such examples, may provide up to a 120-degree horizontal and 35-degree vertical field-of-view, with a 200 m range even for low-reflectivity objects. Front-mounted LIDAR sensor(s) **464** may be configured for a horizontal field of view between 45 degrees and 135 degrees.

(108) In some examples, LIDAR technologies, such as 3D flash LIDAR, may also be used. 3D Flash LIDAR uses a flash of a laser as a transmission source, to illuminate vehicle surroundings up to approximately 200 m. A flash LIDAR unit includes a receptor, which records the laser pulse transit time and the reflected light on each pixel, which in turn corresponds to the range from the vehicle to the objects. Flash LIDAR may allow for highly accurate and distortion-free images of the surroundings to be generated with every laser flash. In some examples, four flash LIDAR sensors may be deployed, one at each side of the vehicle **400**. Available 3D flash LIDAR systems include a solid-state 3D staring array LIDAR camera with no moving parts other than a fan (e.g., a non-scanning LIDAR device). The flash LIDAR device may use a five nanosecond class I (eye-safe) laser pulse per frame and may capture the reflected laser light in the form of 3D range point clouds and co-registered intensity data. By using flash LIDAR, and because flash LIDAR is a solid-state device with no moving parts, the LIDAR sensor(s) **464** may be less susceptible to motion blur, vibration, and/or shock.

(109) The vehicle may further include IMU sensor(s) **466**. The IMU sensor(s) **466** may be located at a center of the rear axle of the vehicle **400**, in some examples. The IMU sensor(s) **466** may include, for example and without limitation, an accelerometer(s), a magnetometer(s), a gyroscope(s), a magnetic compass(es), and/or other sensor types. In some examples, such as in six-axis applications, the IMU sensor(s) **466** may include accelerometers and gyroscopes, while in nine-axis applications, the IMU sensor(s) **466** may include accelerometers, gyroscopes, and magnetometers.

(110) In some embodiments, the IMU sensor(s) **466** may be implemented as a miniature, high-performance GPS-Aided Inertial Navigation System (GPS/INS) that combines micro-electro-mechanical systems (MEMS) inertial sensors, a high-sensitivity GPS receiver, and advanced Kalman filtering algorithms to provide estimates of position, velocity, and attitude. As such, in some examples, the IMU sensor(s) **466** may enable the vehicle **400** to estimate heading without requiring input from a magnetic sensor by directly observing and correlating the changes in velocity from GPS to the IMU sensor(s) **466**. In some examples, the IMU sensor(s) **466** and the GNSS sensor(s) **458** may be combined in a single integrated unit.

(111) The vehicle may include microphone(s) **496** placed in and/or around the vehicle **400**. The microphone(s) **496** may be used for emergency vehicle detection and identification, among other things.

(112) The vehicle may further include any number of camera types, including stereo camera(s) **468**, wide-view camera(s) **470**, infrared camera(s) **472**, surround camera(s) **474**, long-range and/or mid-range camera(s) **498**, and/or other camera types. The cameras may be used to capture image data around an entire periphery of the vehicle **400**. The types of cameras used depends on the embodiments and requirements for the vehicle **400**, and any combination of camera types may be used to provide the necessary coverage around the vehicle **400**. In addition, the number of cameras may differ depending on the embodiment. For example, the vehicle may include six cameras, seven cameras, ten cameras, twelve cameras, and/or another number of cameras. The cameras may support, as an example and without limitation, Gigabit Multimedia Serial Link (GMSL) and/or Gigabit Ethernet. Each of the camera(s) is described with more detail herein with respect to FIG. **4A** and FIG. **4B**.

(113) The vehicle **400** may further include vibration sensor(s) **442**. The vibration sensor(s) **442** may measure vibrations of components of the vehicle, such as the axle(s). For example, changes in vibrations may indicate a change in road surfaces. In another example, when two or more vibration sensors **442** are used, the differences between the vibrations may be used to determine friction or slippage of the road surface (e.g., when the difference in vibration is between a power-driven axle and a freely rotating axle).

(114) The vehicle **400** may include an ADAS system **438**. The ADAS system **438** may include an SoC, in some examples. The ADAS system **438** may include autonomous/adaptive/automatic cruise control (ACC), cooperative adaptive cruise control (CACC), forward crash warning (FCW), automatic emergency braking (AEB), lane departure warnings (LDW), lane keep assist (LKA), blind spot warning (BSW), rear cross-traffic warning (RCTW), collision warning systems (CWS), lane centering (LC), and/or other features and functionality.

(115) The ACC systems may use RADAR sensor(s) **460**, LIDAR sensor(s) **464**, and/or a camera(s). The ACC systems may include longitudinal ACC and/or lateral ACC. Longitudinal ACC monitors and controls the distance to the vehicle immediately ahead of the vehicle **400** and automatically adjusts the vehicle speed to maintain a safe distance from vehicles ahead. Lateral ACC performs distance keeping, and advises the vehicle **400** to change lanes when necessary. Lateral ACC is related to other ADAS applications such as LC and CWS.

(116) CACC uses information from other vehicles that may be received via the network interface **424** and/or the wireless antenna(s) **426** from other vehicles via a wireless link, or indirectly, over a network connection (e.g., over the Internet). Direct links may be provided by a vehicle-to-vehicle (V2V) communication link, while indirect links may be infrastructure-to-vehicle (I2V) communication links. In general, the V2V communication concept provides information about the immediately preceding vehicles (e.g., vehicles immediately ahead of and in the same lane as the vehicle **400**), while the I2V communication concept provides information about traffic farther ahead. CACC systems may include either or both I2V and V2V information sources. Given the information of the vehicles ahead of the vehicle **400**, CACC may be more reliable, and it has potential to improve traffic flow smoothness and reduce congestion on the road.

(117) FCW systems are designed to alert the driver to a hazard, so that the driver may take corrective action. FCW systems use a front-facing camera and/or RADAR sensor(s) **460**, coupled to a dedicated processor, DSP, FPGA, and/or ASIC, that is electrically coupled to driver feedback, such as a display, speaker, and/or vibrating component. FCW systems may provide a warning, such as in the form of a sound, visual warning, vibration and/or a quick brake pulse.

(118) AEB systems detect an impending forward collision with another vehicle or other object, and may automatically apply the brakes if the driver does not take corrective action within a specified time or distance parameter. AEB systems may use front-facing camera(s) and/or RADAR sensor(s) **460**, coupled to a dedicated processor, DSP, FPGA, and/or ASIC. When the AEB system detects a hazard, it typically first alerts the driver to take corrective action to avoid the collision and, if the driver does not take corrective action, the AEB system may automatically apply the brakes in an effort to prevent, or at least mitigate, the impact of the predicted collision. AEB systems, may include techniques such as dynamic brake support and/or crash imminent braking.

(119) LDW systems provide visual, audible, and/or tactile warnings, such as steering wheel or seat vibrations, to alert the driver when the vehicle **400** crosses lane markings. An LDW system does not activate when the driver indicates an intentional lane departure, by activating a turn signal. LDW systems may use front-side facing cameras, coupled to a dedicated processor, DSP, FPGA, and/or ASIC, that is electrically coupled to driver feedback, such as a display, speaker, and/or vibrating component.

(120) LKA systems are a variation of LDW systems. LKA systems provide steering input or braking to correct the vehicle **400** if the vehicle **400** starts to exit the lane. BSW systems detect and warn the driver of vehicles in an automobile's blind spot. BSW systems may provide a visual,

audible, and/or tactile alert to indicate that merging or changing lanes is unsafe. The system may provide an additional warning when the driver uses a turn signal. BSW systems may use rear-side facing camera(s) and/or RADAR sensor(s) **460**, coupled to a dedicated processor, DSP, FPGA, and/or ASIC, that is electrically coupled to driver feedback, such as a display, speaker, and/or vibrating component.

(121) RCTW systems may provide visual, audible, and/or tactile notification when an object is detected outside the rear-camera range when the vehicle **400** is backing up. Some RCTW systems include AEB to ensure that the vehicle brakes are applied to avoid a crash. RCTW systems may use one or more rear-facing RADAR sensor(s) **460**, coupled to a dedicated processor, DSP, FPGA, and/or ASIC, that is electrically coupled to driver feedback, such as a display, speaker, and/or vibrating component.

(122) Conventional ADAS systems may be prone to false positive results, which may be annoying and distracting to a driver, but typically are not catastrophic, because the ADAS systems alert the driver and allow the driver to decide whether a safety condition truly exists and act accordingly. However, in an autonomous vehicle **400**, the vehicle **400** itself must, in the case of conflicting results, decide whether to heed the result from a primary computer or a secondary computer (e.g., a first controller **436** or a second controller **436**). For example, in some embodiments, the ADAS system **438** may be a backup and/or secondary computer for providing perception information to a backup computer rationality module. The backup computer rationality monitor may run a redundant diverse software on hardware components to detect faults in perception and dynamic driving tasks. Outputs from the ADAS system **438** may be provided to a supervisory MCU. If outputs from the primary computer and the secondary computer conflict, the supervisory MCU must determine how to reconcile the conflict to ensure safe operation.

(123) In some examples, the primary computer may be configured to provide the supervisory MCU with a confidence score, indicating the primary computer's confidence in the chosen result. If the confidence score exceeds a threshold, the supervisory MCU may follow the primary computer's direction, regardless of whether the secondary computer provides a conflicting or inconsistent result. Where the confidence score does not meet the threshold, and where the primary and secondary computer indicate different results (e.g., the conflict), the supervisory MCU may arbitrate between the computers to determine the appropriate outcome.

(124) The supervisory MCU may be configured to run a neural network(s) that is trained and configured to determine, based on outputs from the primary computer and the secondary computer, conditions under which the secondary computer provides false alarms. Thus, the neural network(s) in the supervisory MCU may learn when the secondary computer's output can be trusted, and when it cannot. For example, when the secondary computer is a RADAR-based FCW system, a neural network(s) in the supervisory MCU may learn when the FCW system is identifying metallic objects that are not, in fact, hazards, such as a drainage grate or manhole cover that triggers an alarm. Similarly, when the secondary computer is a camera-based LDW system, a neural network in the supervisory MCU may learn to override the LDW when bicyclists or pedestrians are present and a lane departure is, in fact, the safest maneuver. In embodiments that include a neural network(s) running on the supervisory MCU, the supervisory MCU may include at least one of a DLA or GPU suitable for running the neural network(s) with associated memory. In preferred embodiments, the supervisory MCU may comprise and/or be included as a component of the SoC(s) **404**.

(125) In other examples, ADAS system **438** may include a secondary computer that performs ADAS functionality using traditional rules of computer vision. As such, the secondary computer may use classic computer vision rules (if-then), and the presence of a neural network(s) in the supervisory MCU may improve reliability, safety and performance. For example, the diverse implementation and intentional non-identity make the overall system more fault-tolerant, especially to faults caused by software (or software-hardware interface) functionality. For example, if there is a software bug or error in the software running on the primary computer, and the non-identical

software code running on the secondary computer provides the same overall result, the supervisory MCU may have greater confidence that the overall result is correct, and the bug in software or hardware used by the primary computer is not causing material error.

(126) In some examples, the output of the ADAS system **438** may be fed into the primary computer's perception block and/or the primary computer's dynamic driving task block. For example, if the ADAS system **438** indicates a forward crash warning due to an object immediately ahead, the perception block may use this information when identifying objects. In other examples, the secondary computer may have its own neural network that is trained and thus reduces the risk of false positives, as described herein.

(127) The vehicle **400** may further include the infotainment SoC **430** (e.g., an in-vehicle infotainment system (IVI)). Although illustrated and described as an SoC, the infotainment system may not be a SoC, and may include two or more discrete components. The infotainment SoC **430** may include a combination of hardware and software that may be used to provide audio (e.g., music, a personal digital assistant, navigational instructions, news, radio, etc.), video (e.g., TV, movies, streaming, etc.), phone (e.g., hands-free calling), network connectivity (e.g., LTE, Wi-Fi, etc.), and/or information services (e.g., navigation systems, rear-parking assistance, a radio data system, vehicle-related information such as fuel level, total distance covered, brake fuel level, oil level, door open/close, air filter information, etc.) to the vehicle **400**. For example, the infotainment SoC **430** may include radios, disk players, navigation systems, video players, USB and Bluetooth connectivity, carputers, in-car entertainment, Wi-Fi, steering wheel audio controls, hands-free voice control, a heads-up display (HUD), an HMI display **434**, a telematics device, a control panel (e.g., for controlling and/or interacting with various components, features, and/or systems), and/or other components. The infotainment SoC **430** may further be used to provide information (e.g., visual and/or audible) to a user(s) of the vehicle, such as information from the ADAS system **438**, autonomous driving information such as planned vehicle maneuvers, trajectories, surrounding environment information (e.g., intersection information, vehicle information, road information, etc.), and/or other information.

(128) The infotainment SoC **430** may include GPU functionality. The infotainment SoC **430** may communicate over the bus **402** (e.g., CAN bus, Ethernet, etc.) with other devices, systems, and/or components of the vehicle **400**. In some examples, the infotainment SoC **430** may be coupled to a supervisory MCU such that the GPU of the infotainment system may perform some self-driving functions in the event that the primary controller(s) **436** (e.g., the primary and/or backup computers of the vehicle **400**) fail. In such an example, the infotainment SoC **430** may put the vehicle **400** into a chauffeur to safe-stop mode, as described herein.

(129) The vehicle **400** may further include an instrument cluster **432** (e.g., a digital dash, an electronic instrument cluster, a digital instrument panel, etc.). The instrument cluster **432** may include a controller and/or supercomputer (e.g., a discrete controller or supercomputer). The instrument cluster **432** may include a set of instrumentation such as a speedometer, fuel level, oil pressure, tachometer, odometer, turn indicators, gearshift position indicator, seat belt warning light(s), parking-brake warning light(s), engine-malfunction light(s), airbag (SRS) system information, lighting controls, safety system controls, navigation information, etc. In some examples, information may be displayed and/or shared among the infotainment SoC **430** and the instrument cluster **432**. In other words, the instrument cluster **432** may be included as part of the infotainment SoC **430**, or vice versa.

(130) FIG. **4D** is a system diagram for communication between cloud-based server(s) and the example autonomous vehicle **400** of FIG. **4A**, in accordance with some embodiments of the present disclosure. The system **476** may include server(s) **478**, network(s) **490**, and vehicles, including the vehicle **400**. The server(s) **478** may include a plurality of GPUs **484(A)-484(H)** (collectively referred to herein as GPUs **484**), PCIe switches **482(A)-482(H)** (collectively referred to herein as PCIe switches **482**), and/or CPUs **480(A)-480(B)** (collectively referred to herein as CPUs **480**).

The GPUs **484**, the CPUs **480**, and the PCIe switches may be interconnected with high-speed interconnects such as, for example and without limitation, NVLink interfaces **488** developed by NVIDIA and/or PCIe connections **486**. In some examples, the GPUs **484** are connected via NVLink and/or NVSwitch SoC and the GPUs **484** and the PCIe switches **482** are connected via PCIe interconnects. Although eight GPUs **484**, two CPUs **480**, and two PCIe switches are illustrated, this is not intended to be limiting. Depending on the embodiment, each of the server(s) **478** may include any number of GPUs **484**, CPUs **480**, and/or PCIe switches. For example, the server(s) **478** may each include eight, sixteen, thirty-two, and/or more GPUs **484**.

(131) The server(s) **478** may receive, over the network(s) **490** and from the vehicles, image data representative of images showing unexpected or changed road conditions, such as recently commenced roadwork. The server(s) **478** may transmit, over the network(s) **490** and to the vehicles, neural networks **492**, updated neural networks **492**, and/or map information **494**, including information regarding traffic and road conditions. The updates to the map information **494** may include updates for the HD map **422**, such as information regarding construction sites, potholes, detours, flooding, and/or other obstructions. In some examples, the neural networks **492**, the updated neural networks **492**, and/or the map information **494** may have resulted from new training and/or experiences represented in data received from any number of vehicles in the environment, and/or based on training performed at a datacenter (e.g., using the server(s) **478** and/or other servers).

(132) The server(s) **478** may be used to train machine learning models (e.g., neural networks) based on training data. The training data may be generated by the vehicles, and/or may be generated in a simulation (e.g., using a game engine). In some examples, the training data is tagged (e.g., where the neural network benefits from supervised learning) and/or undergoes other pre-processing, while in other examples the training data is not tagged and/or pre-processed (e.g., where the neural network does not require supervised learning). Training may be executed according to any one or more classes of machine learning techniques, including, without limitation, classes such as: supervised training, semi-supervised training, unsupervised training, self-learning, reinforcement learning, federated learning, transfer learning, feature learning (including principal component and cluster analyses), multi-linear subspace learning, manifold learning, representation learning (including sparse dictionary learning), rule-based machine learning, anomaly detection, and any variants or combinations therefor. Once the machine learning models are trained, the machine learning models may be used by the vehicles (e.g., transmitted to the vehicles over the network(s) **490**, and/or the machine learning models may be used by the server(s) **478** to remotely monitor the vehicles.

(133) In some examples, the server(s) **478** may receive data from the vehicles and apply the data to up-to-date real-time neural networks for real-time inferencing. The server(s) **478** may include deep-learning supercomputers and/or dedicated AI computers powered by GPU(s) **484**, such as a DGX and DGX Station machines developed by NVIDIA. However, in some examples, the server(s) **478** may include deep learning infrastructure that use only CPU-powered datacenters.

(134) The deep-learning infrastructure of the server(s) **478** may be capable of fast, real-time inferencing, and may use that capability to evaluate and verify the health of the processors, software, and/or associated hardware in the vehicle **400**. For example, the deep-learning infrastructure may receive periodic updates from the vehicle **400**, such as a sequence of images and/or objects that the vehicle **400** has located in that sequence of images (e.g., via computer vision and/or other machine learning object classification techniques). The deep-learning infrastructure may run its own neural network to identify the objects and compare them with the objects identified by the vehicle **400** and, if the results do not match and the infrastructure concludes that the AI in the vehicle **400** is malfunctioning, the server(s) **478** may transmit a signal to the vehicle **400** instructing a fail-safe computer of the vehicle **400** to assume control, notify the passengers, and complete a safe parking maneuver.

(135) For inferencing, the server(s) **478** may include the GPU(s) **484** and one or more programmable inference accelerators (e.g., NVIDIA's TensorRT). The combination of GPU-powered servers and inference acceleration may make real-time responsiveness possible. In other examples, such as where performance is less critical, servers powered by CPUs, FPGAs, and other processors may be used for inferencing.

(136) FIG. 5 is a block diagram of an example computing device(s) **500** suitable for use in implementing some embodiments of the present disclosure. Computing device **500** may include an interconnect system **502** that directly or indirectly couples the following devices: memory **504**, one or more central processing units (CPUs) **506**, one or more graphics processing units (GPUs) **508**, a communication interface **510**, I/O ports **512**, input/output components **514**, a power supply **516**, one or more presentation components **518** (e.g., display(s)), and one or more logic units **520**.

(137) Although the various blocks of FIG. 5 are shown as connected via the interconnect system **502** with lines, this is not intended to be limiting and is for clarity only. For example, in some embodiments, a presentation component **518**, such as a display device, may be considered an I/O component **514** (e.g., if the display is a touch screen). As another example, the CPUs **506** and/or GPUs **508** may include memory (e.g., the memory **504** may be representative of a storage device in addition to the memory of the GPUs **508**, the CPUs **506**, and/or other components). In other words, the computing device of FIG. 5 is merely illustrative. Distinction is not made between such categories as “workstation,” “server,” “laptop,” “desktop,” “tablet,” “client device,” “mobile device,” “hand-held device,” “game console,” “electronic control unit (ECU),” “virtual reality system,” “augmented reality system,” and/or other device or system types, as all are contemplated within the scope of the computing device of FIG. 5.

(138) The interconnect system **502** may represent one or more links or busses, such as an address bus, a data bus, a control bus, or a combination thereof. The interconnect system **502** may include one or more bus or link types, such as an industry standard architecture (ISA) bus, an extended industry standard architecture (EISA) bus, a video electronics standards association (VESA) bus, a peripheral component interconnect (PCI) bus, a peripheral component interconnect express (PCIe) bus, and/or another type of bus or link. In some embodiments, there are direct connections between components. As an example, the CPU **506** may be directly connected to the memory **504**. Further, the CPU **506** may be directly connected to the GPU **508**. Where there is direct, or point-to-point, connection between components, the interconnect system **502** may include a PCIe link to carry out the connection. In these examples, a PCI bus need not be included in the computing device **500**.

(139) The memory **504** may include any of a variety of computer-readable media. The computer-readable media may be any available media that may be accessed by the computing device **500**. The computer-readable media may include both volatile and nonvolatile media, and removable and non-removable media. By way of example, and not limitation, the computer-readable media may comprise computer-storage media and communication media.

(140) The computer-storage media may include both volatile and nonvolatile media and/or removable and non-removable media implemented in any method or technology for storage of information such as computer-readable instructions, data structures, program modules, and/or other data types. For example, the memory **504** may store computer-readable instructions (e.g., that represent a program(s) and/or a program element(s), such as an operating system. Computer-storage media may include, but is not limited to, RAM, ROM, EEPROM, flash memory or other memory technology, CD-ROM, digital versatile disks (DVD) or other optical disk storage, magnetic cassettes, magnetic tape, magnetic disk storage or other magnetic storage devices, or any other medium that may be used to store the desired information and that may be accessed by computing device **500**. As used herein, computer storage media does not comprise signals per se.

(141) The computer storage media may embody computer-readable instructions, data structures, program modules, and/or other data types in a modulated data signal such as a carrier wave or other transport mechanism and includes any information delivery media. The term “modulated data

signal” may refer to a signal that has one or more of its characteristics set or changed in such a manner as to encode information in the signal. By way of example, and not limitation, the computer storage media may include wired media such as a wired network or direct-wired connection, and wireless media such as acoustic, RF, infrared and other wireless media. Combinations of any of the above should also be included within the scope of computer-readable media.

(142) The CPU(s) **506** may be configured to execute at least some of the computer-readable instructions to control one or more components of the computing device **500** to perform one or more of the methods and/or processes described herein. The CPU(s) **506** may each include one or more cores (e.g., one, two, four, eight, twenty-eight, seventy-two, etc.) that are capable of handling a multitude of software threads simultaneously. The CPU(s) **506** may include any type of processor, and may include different types of processors depending on the type of computing device **500** implemented (e.g., processors with fewer cores for mobile devices and processors with more cores for servers). For example, depending on the type of computing device **500**, the processor may be an Advanced RISC Machines (ARM) processor implemented using Reduced Instruction Set Computing (RISC) or an x86 processor implemented using Complex Instruction Set Computing (CISC). The computing device **500** may include one or more CPUs **506** in addition to one or more microprocessors or supplementary co-processors, such as math co-processors.

(143) In addition to or alternatively from the CPU(s) **506**, the GPU(s) **508** may be configured to execute at least some of the computer-readable instructions to control one or more components of the computing device **500** to perform one or more of the methods and/or processes described herein. One or more of the GPU(s) **508** may be an integrated GPU (e.g., with one or more of the CPU(s) **506** and/or one or more of the GPU(s) **508** may be a discrete GPU. In embodiments, one or more of the GPU(s) **508** may be a coprocessor of one or more of the CPU(s) **506**. The GPU(s) **508** may be used by the computing device **500** to render graphics (e.g., 3D graphics) or perform general purpose computations. For example, the GPU(s) **508** may be used for General-Purpose computing on GPUs (GPGPU). The GPU(s) **508** may include hundreds or thousands of cores that are capable of handling hundreds or thousands of software threads simultaneously. The GPU(s) **508** may generate pixel data for output images in response to rendering commands (e.g., rendering commands from the CPU(s) **506** received via a host interface). The GPU(s) **508** may include graphics memory, such as display memory, for storing pixel data or any other suitable data, such as GPGPU data. The display memory may be included as part of the memory **504**. The GPU(s) **508** may include two or more GPUs operating in parallel (e.g., via a link). The link may directly connect the GPUs (e.g., using NVLINK) or may connect the GPUs through a switch (e.g., using NVSwitch). When combined together, each GPU **508** may generate pixel data or GPGPU data for different portions of an output or for different outputs (e.g., a first GPU for a first image and a second GPU for a second image). Each GPU may include its own memory, or may share memory with other GPUs.

(144) In addition to or alternatively from the CPU(s) **506** and/or the GPU(s) **508**, the logic unit(s) **520** may be configured to execute at least some of the computer-readable instructions to control one or more components of the computing device **500** to perform one or more of the methods and/or processes described herein. In embodiments, the CPU(s) **506**, the GPU(s) **508**, and/or the logic unit(s) **520** may discretely or jointly perform any combination of the methods, processes and/or portions thereof. One or more of the logic units **520** may be part of and/or integrated in one or more of the CPU(s) **506** and/or the GPU(s) **508** and/or one or more of the logic units **520** may be discrete components or otherwise external to the CPU(s) **506** and/or the GPU(s) **508**. In embodiments, one or more of the logic units **520** may be a coprocessor of one or more of the CPU(s) **506** and/or one or more of the GPU(s) **508**.

(145) Examples of the logic unit(s) **520** include one or more processing cores and/or components thereof, such as Tensor Cores (TCs), Tensor Processing Units (TPUs), Pixel Visual Cores (PVCs),

Vision Processing Units (VPUs), Graphics Processing Clusters (GPCs), Texture Processing Clusters (TPCs), Streaming Multiprocessors (SMs), Tree Traversal Units (TTUs), Artificial Intelligence Accelerators (AIAs), Deep Learning Accelerators (DLAs), Arithmetic-Logic Units (ALUs), Application-Specific Integrated Circuits (ASICs), Floating Point Units (FPUs), I/O elements, peripheral component interconnect (PCI) or peripheral component interconnect express (PCIe) elements, and/or the like.

(146) The communication interface **510** may include one or more receivers, transmitters, and/or transceivers that enable the computing device **500** to communicate with other computing devices via an electronic communication network, including wired and/or wireless communications. The communication interface **510** may include components and functionality to enable communication over any of a number of different networks, such as wireless networks (e.g., Wi-Fi, Z-Wave, Bluetooth, Bluetooth LE, ZigBee, etc.), wired networks (e.g., communicating over Ethernet or InfiniBand), low-power wide-area networks (e.g., LoRaWAN, SigFox, etc.), and/or the Internet.

(147) The I/O ports **512** may enable the computing device **500** to be logically coupled to other devices including the I/O components **514**, the presentation component(s) **518**, and/or other components, some of which may be built into (e.g., integrated in) the computing device **500**. Illustrative I/O components **514** include a microphone, mouse, keyboard, joystick, game pad, game controller, satellite dish, scanner, printer, wireless device, etc. The I/O components **514** may provide a natural user interface (NUI) that processes air gestures, voice, or other physiological inputs generated by a user. In some instances, inputs may be transmitted to an appropriate network element for further processing. An NUI may implement any combination of speech recognition, stylus recognition, facial recognition, biometric recognition, gesture recognition both on screen and adjacent to the screen, air gestures, head and eye tracking, and touch recognition (as described in more detail below) associated with a display of the computing device **500**. The computing device **500** may include depth cameras, such as stereoscopic camera systems, infrared camera systems, RGB camera systems, touchscreen technology, and combinations of these, for gesture detection and recognition. Additionally, the computing device **500** may include accelerometers or gyroscopes (e.g., as part of an inertia measurement unit (IMU)) that enable detection of motion. In some examples, the output of the accelerometers or gyroscopes may be used by the computing device **500** to render immersive augmented reality or virtual reality.

(148) The power supply **516** may include a hard-wired power supply, a battery power supply, or a combination thereof. The power supply **516** may provide power to the computing device **500** to enable the components of the computing device **500** to operate.

(149) The presentation component(s) **518** may include a display (e.g., a monitor, a touch screen, a television screen, a heads-up-display (HUD), other display types, or a combination thereof), speakers, and/or other presentation components. The presentation component(s) **518** may receive data from other components (e.g., the GPU(s) **508**, the CPU(s) **506**, etc.), and output the data (e.g., as an image, video, sound, etc.).

(150) The disclosure may be described in the general context of computer code or machine-useable instructions, including computer-executable instructions such as program modules, being executed by a computer or other machine, such as a personal data assistant or other handheld device. Generally, program modules including routines, programs, objects, components, data structures, etc., refer to codes that perform particular tasks or implement particular abstract data types. The disclosure may be practiced in a variety of system configurations, including hand-held devices, consumer electronics, general-purpose computers, more specialty computing devices, etc. The disclosure may also be practiced in distributed computing environments where tasks are performed by remote-processing devices that are linked through a communications network.

(151) FIG. **6A** illustrates an exemplary architecture in which a plurality of GPUs **610(1)-610(N)** is communicatively coupled to a plurality of multi-core processors **605(1)-605(M)** over high-speed links **640(1)-640(N)** (e.g., buses, point-to-point interconnects, etc.). In at least one embodiment,

high-speed links **640(1)-640(N)** support a communication throughput of 4 GB/s, 30 GB/s, 80 GB/s or higher. In at least one embodiment, various interconnect protocols may be used including, but not limited to, PCIe 4.0 or 5.0 and NVLink 2.0. In various figures, “N” and “M” represent positive integers, values of which may be different from figure to figure.

(152) In addition, and in at least one embodiment, two or more of GPUs **610** are interconnected over high-speed links **629(1)-629(2)**, which may be implemented using similar or different protocols/links than those used for high-speed links **640(1)-640(N)**. Similarly, two or more of multi-core processors **605** may be connected over a high-speed link **628** which may be symmetric multi-processor (SMP) buses operating at 20 GB/s, 30 GB/s, 120 GB/s or higher. Alternatively, all communication between various system components shown in FIG. **6A** may be accomplished using similar protocols/links (e.g., over a common interconnection fabric).

(153) In at least one embodiment, each multi-core processor **605** is communicatively coupled to a processor memory **601(1)-601(M)**, via memory interconnects **626(1)-626(M)**, respectively, and each GPU **610(1)-610(N)** is communicatively coupled to GPU memory **620(1)-620(N)** over GPU memory interconnects **650(1)-650(N)**, respectively. In at least one embodiment, memory interconnects **626** and **650** may utilize similar or different memory access technologies. By way of example, and not limitation, processor memories **601(1)-601(M)** and GPU memories **620** may be volatile memories such as dynamic random access memories (DRAMs) (including stacked DRAMs), Graphics DDR SDRAM (GDDR) (e.g., GDDR5, GDDR6), or High Bandwidth Memory (HBM) and/or may be non-volatile memories such as 3D XPoint or Nano-Ram. In at least one embodiment, some portion of processor memories **601** may be volatile memory and another portion may be non-volatile memory (e.g., using a two-level memory (2LM) hierarchy).

(154) As described herein, although various multi-core processors **605** and GPUs **610** may be physically coupled to a particular memory **601**, **620**, respectively, and/or a unified memory architecture may be implemented in which a virtual system address space (also referred to as “effective address” space) is distributed among various physical memories. For example, processor memories **601(1)-601(M)** may each comprise 64 GB of system memory address space and GPU memories **620(1)-620(N)** may each comprise 32 GB of system memory address space resulting in a total of 256 GB addressable memory when M=2 and N=4. Other values for N and M are possible.

(155) FIG. **6B** illustrates additional details for an interconnection between a multi-core processor **607** and a graphics acceleration module **646** in accordance with one exemplary embodiment. In at least one embodiment, graphics acceleration module **646** may include one or more GPU chips integrated on a line card which is coupled to processor **607** via high-speed link **640** (e.g., a PCIe bus, NVLink, etc.). In at least one embodiment, graphics acceleration module **646** may alternatively be integrated on a package or chip with processor **607**.

(156) In at least one embodiment, processor **607** includes a plurality of cores **660A-660D**, each with a translation lookaside buffer (“TLB”) **661A-661D** and one or more caches **662A-662D**. In at least one embodiment, cores **660A-660D** may include various other components for executing instructions and processing data that are not illustrated. In at least one embodiment, caches **662A-662D** may comprise Level 1 (L1) and Level 2 (L2) caches. In addition, one or more shared caches **656** may be included in caches **662A-662D** and shared by sets of cores **660A-660D**. For example, one embodiment of processor **607** includes 24 cores, each with its own L1 cache, twelve shared L2 caches, and twelve shared L3 caches. In this embodiment, one or more L2 and L3 caches are shared by two adjacent cores. In at least one embodiment, processor **607** and graphics acceleration module **646** connect with system memory **614**, which may include processor memories **601(1)-601(M)** of FIG. **6A**.

(157) In at least one embodiment, coherency is maintained for data and instructions stored in various caches **662A-662D**, **656** and system memory **614** via inter-core communication over a coherence bus **664**. In at least one embodiment, for example, each cache may have cache coherency logic/circuitry associated therewith to communicate to over coherence bus **664** in response to

detected reads or writes to particular cache lines. In at least one embodiment, a cache snooping protocol is implemented over coherence bus **664** to snoop cache accesses.

(158) In at least one embodiment, a proxy circuit **625** communicatively couples graphics acceleration module **646** to coherence bus **664**, allowing graphics acceleration module **646** to participate in a cache coherence protocol as a peer of cores **660A-660D**. In particular, in at least one embodiment, an interface **635** provides connectivity to proxy circuit **625** over high-speed link **640** and an interface **637** connects graphics acceleration module **646** to high-speed link **640**.

(159) In at least one embodiment, an accelerator integration circuit **636** provides cache management, memory access, context management, and interrupt management services on behalf of a plurality of graphics processing engines **631(1)-631(N)** of graphics acceleration module **646**. In at least one embodiment, graphics processing engines **631(1)-631(N)** may each comprise a separate graphics processing unit (GPU). In at least one embodiment, graphics processing engines **631(1)-631(N)** alternatively may comprise different types of graphics processing engines within a GPU, such as graphics execution units, media processing engines (e.g., video encoders/decoders), samplers, and blit engines. In at least one embodiment, graphics acceleration module **646** may be a GPU with a plurality of graphics processing engines **631(1)-631(N)** or graphics processing engines **631(1)-631(N)** may be individual GPUs integrated on a common package, line card, or chip.

(160) In at least one embodiment, accelerator integration circuit **636** includes a memory management unit (MMU) **639** for performing various memory management functions such as virtual-to-physical memory translations (also referred to as effective-to-real memory translations) and memory access protocols for accessing system memory **614**. In at least one embodiment, MMU **639** may also include a translation lookaside buffer (TLB) (not shown) for caching virtual/effective to physical/real address translations. In at least one embodiment, a cache **638** can store commands and data for efficient access by graphics processing engines **631(1)-631(N)**. In at least one embodiment, data stored in cache **638** and graphics memories **633(1)-633(M)** is kept coherent with core caches **662A-662D**, **656** and system memory **614**, possibly using a fetch unit **644**. As mentioned, this may be accomplished via proxy circuit **625** on behalf of cache **638** and memories **633(1)-633(M)** (e.g., sending updates to cache **638** related to modifications/accesses of cache lines on processor caches **662A-662D**, **656** and receiving updates from cache **638**).

(161) In at least one embodiment, a set of registers **645** store context data for threads executed by graphics processing engines **631(1)-631(N)** and a context management circuit **648** manages thread contexts. For example, context management circuit **648** may perform save and restore operations to save and restore contexts of various threads during contexts switches (e.g., where a first thread is saved and a second thread is stored so that a second thread can be execute by a graphics processing engine). For example, on a context switch, context management circuit **648** may store current register values to a designated region in memory (e.g., identified by a context pointer). It may then restore register values when returning to a context. In at least one embodiment, an interrupt management circuit **647** receives and processes interrupts received from system devices.

(162) In at least one embodiment, virtual/effective addresses from a graphics processing engine **631** are translated to real/physical addresses in system memory **614** by MMU **639**. In at least one embodiment, accelerator integration circuit **636** supports multiple (e.g., 4, 8, 16) graphics accelerator modules **646** and/or other accelerator devices. In at least one embodiment, graphics accelerator module **646** may be dedicated to a single application executed on processor **607** or may be shared between multiple applications. In at least one embodiment, a virtualized graphics execution environment is presented in which resources of graphics processing engines **631(1)-631(N)** are shared with multiple applications or virtual machines (VMs). In at least one embodiment, resources may be subdivided into “slices” which are allocated to different VMs and/or applications based on processing requirements and priorities associated with VMs and/or applications.

(163) In at least one embodiment, accelerator integration circuit **636** performs as a bridge to a system for graphics acceleration module **646** and provides address translation and system memory

cache services. In addition, in at least one embodiment, accelerator integration circuit **636** may provide virtualization facilities for a host processor to manage virtualization of graphics processing engines **631(1)-63 l(N)**, interrupts, and memory management.

(164) In at least one embodiment, because hardware resources of graphics processing engines **631(1)-63 l(N)** are mapped explicitly to a real address space seen by host processor **607**, any host processor can address these resources directly using an effective address value. In at least one embodiment, one function of accelerator integration circuit **636** is physical separation of graphics processing engines **631(1)-63 l(N)** so that they appear to a system as independent units.

(165) In at least one embodiment, one or more graphics memories **633(1)-633(M)** are coupled to each of graphics processing engines **631(1)-63 l(N)**, respectively and $N=M$. In at least one embodiment, graphics memories **633(1)-633(M)** store instructions and data being processed by each of graphics processing engines **631(1)-63 l(N)**. In at least one embodiment, graphics memories **633(1)-633(M)** may be volatile memories such as DRAMs (including stacked DRAMs), GDDR memory (e.g., GDDR5, GDDR6), or HBM, and/or may be non-volatile memories such as 3D XPoint or Nano-Ram.

(166) In at least one embodiment, to reduce data traffic over high-speed link **640**, biasing techniques can be used to ensure that data stored in graphics memories **633(1)-633(M)** is data that will be used most frequently by graphics processing engines **631(1)-63 l(N)** and preferably not used by cores **660A-660D** (at least not frequently). Similarly, in at least one embodiment, a biasing mechanism attempts to keep data needed by cores (and preferably not graphics processing engines **631(1)-63 l(N)**) within caches **662A-662D**, **656** and system memory **614**.

(167) FIG. **6C** illustrates another exemplary embodiment in which accelerator integration circuit **636** is integrated within processor **607**. In this embodiment, graphics processing engines **631(1)-63 l(N)** communicate directly over high-speed link **640** to accelerator integration circuit **636** via interface **637** and interface **635** (which, again, may be any form of bus or interface protocol). In at least one embodiment, accelerator integration circuit **636** may perform similar operations as those described with respect to FIG. **6B**, but potentially at a higher throughput given its close proximity to coherence bus **664** and caches **662A-662D**, **656**. In at least one embodiment, an accelerator integration circuit supports different programming models including a dedicated-process programming model (no graphics acceleration module virtualization) and shared programming models (with virtualization), which may include programming models which are controlled by accelerator integration circuit **636** and programming models which are controlled by graphics acceleration module **646**.

(168) In at least one embodiment, graphics processing engines **631(1)-63 l(N)** are dedicated to a single application or process under a single operating system. In at least one embodiment, a single application can funnel other application requests to graphics processing engines **631(1)-63 l(N)**, providing virtualization within a VM/partition.

(169) In at least one embodiment, graphics processing engines **631(1)-63 l(N)**, may be shared by multiple VM/application partitions. In at least one embodiment, shared models may use a system hypervisor to virtualize graphics processing engines **631(1)-63 l(N)** to allow access by each operating system. In at least one embodiment, for single-partition systems without a hypervisor, graphics processing engines **631(1)-63 l(N)** are owned by an operating system. In at least one embodiment, an operating system can virtualize graphics processing engines **631(1)-63 l(N)** to provide access to each process or application.

(170) In at least one embodiment, graphics acceleration module **646** or an individual graphics processing engine **631(1)-63 l(N)** selects a process element using a process handle. In at least one embodiment, process elements are stored in system memory **614** and are addressable using an effective address to real address translation technique described herein. In at least one embodiment, a process handle may be an implementation-specific value provided to a host process when registering its context with graphics processing engine **631(1)-63 l(N)** (that is, calling system

software to add a process element to a process element linked list). In at least one embodiment, a lower 16-bits of a process handle may be an offset of a process element within a process element linked list.

(171) FIG. 6D illustrates an exemplary accelerator integration slice 690. In at least one embodiment, a “slice” comprises a specified portion of processing resources of accelerator integration circuit 636. In at least one embodiment, an application is effective address space 682 within system memory 614 stores process elements 683. In at least one embodiment, process elements 683 are stored in response to GPU invocations 681 from applications 680 executed on processor 607. In at least one embodiment, a process element 683 contains process state for corresponding application 680. In at least one embodiment, a work descriptor (WD) 684 contained in process element 683 can be a single job requested by an application or may contain a pointer to a queue of jobs. In at least one embodiment, WD 684 is a pointer to a job request queue in an application's effective address space 682.

(172) In at least one embodiment, graphics acceleration module 646 and/or individual graphics processing engines 631(1)-631(N) can be shared by all or a subset of processes in a system. In at least one embodiment, an infrastructure for setting up process states and sending a WD 684 to a graphics acceleration module 646 to start a job in a virtualized environment may be included.

(173) In at least one embodiment, a dedicated-process programming model is implementation-specific. In at least one embodiment, in this model, a single process owns graphics acceleration module 646 or an individual graphics processing engine 631. In at least one embodiment, when graphics acceleration module 646 is owned by a single process, a hypervisor initializes accelerator integration circuit 636 for an owning partition and an operating system initializes accelerator integration circuit 636 for an owning process when graphics acceleration module 646 is assigned.

(174) In at least one embodiment, in operation, a WD fetch unit 691 in accelerator integration slice 690 fetches next WD 684, which includes an indication of work to be done by one or more graphics processing engines of graphics acceleration module 646. In at least one embodiment, data from WD 684 may be stored in registers 645 and used by MMU 639, interrupt management circuit 647 and/or context management circuit 648 as illustrated. For example, one embodiment of MMU 639 includes segment/page walk circuitry for accessing segment/page tables 686 within an OS virtual address space 685. In at least one embodiment, interrupt management circuit 647 may process interrupt events 692 received from graphics acceleration module 646. In at least one embodiment, when performing graphics operations, an effective address 693 generated by a graphics processing engine 631(1)-631(N) is translated to a real address by MMU 639.

(175) In at least one embodiment, registers 645 are duplicated for each graphics processing engine 631(1)-631(N) and/or graphics acceleration module 646 and may be initialized by a hypervisor or an operating system. In at least one embodiment, each of these duplicated registers may be included in an accelerator integration slice 690. Exemplary registers that may be initialized by a hypervisor are shown in Table 1.

(176) TABLE-US-00001 TABLE 1 Hypervisor Initialized Registers

Register #	Description
1	Slice Control Register
2	Real Address (RA) Scheduled Processes Area Pointer
3	Authority Mask Override Register
4	Interrupt Vector Table Entry Offset
5	Interrupt Vector Table Entry Limit
6	State Register
7	Logical Partition ID
8	Real address (RA) Hypervisor Accelerator Utilization Record Pointer
9	Storage Description Register

(177) Exemplary registers that may be initialized by an operating system are shown in Table 2.

(178) TABLE-US-00002 TABLE 2 Operating System Initialized Registers

Register #	Description
1	Process and Thread Identification
2	Effective Address (EA) Context Save/Restore Pointer
3	Virtual Address (VA) Accelerator Utilization Record Pointer
4	Virtual Address (VA) Storage Segment Table Pointer
5	Authority Mask
6	Work descriptor

(179) In at least one embodiment, each WD 684 is specific to a particular graphics acceleration module 646 and/or graphics processing engines 631(1)-631(N). In at least one embodiment, it

contains all information required by a graphics processing engine **631(1)-63 l(N)** to do work, or it can be a pointer to a memory location where an application has set up a command queue of work to be completed.

(180) FIG. **6E** illustrates additional details for one exemplary embodiment of a shared model. This embodiment includes a hypervisor real address space **698** in which a process element list **699** is stored. In at least one embodiment, hypervisor real address space **698** is accessible via a hypervisor **696** which virtualizes graphics acceleration module engines for operating system **695**.

(181) In at least one embodiment, shared programming models allow for all or a subset of processes from all or a subset of partitions in a system to use a graphics acceleration module **646**. In at least one embodiment, there are two programming models where graphics acceleration module **646** is shared by multiple processes and partitions, namely time-sliced shared and graphics directed shared.

(182) In at least one embodiment, in this model, system hypervisor **696** owns graphics acceleration module **646** and makes its function available to all operating systems **695**. In at least one embodiment, for a graphics acceleration module **646** to support virtualization by system hypervisor **696**, graphics acceleration module **646** may adhere to certain requirements, such as (1) an application's job request must be autonomous (that is, state does not need to be maintained between jobs), or graphics acceleration module **646** must provide a context save and restore mechanism, (2) an application's job request is guaranteed by graphics acceleration module **646** to complete in a specified amount of time, including any translation faults, or graphics acceleration module **646** provides an ability to preempt processing of a job, and (3) graphics acceleration module **646** must be guaranteed fairness between processes when operating in a directed shared programming model.

(183) In at least one embodiment, application **680** is required to make an operating system **695** system call with a graphics acceleration module type, a work descriptor (WD), an authority mask register (AMR) value, and a context save/restore area pointer (CSRP). In at least one embodiment, graphics acceleration module type describes a targeted acceleration function for a system call. In at least one embodiment, graphics acceleration module type may be a system-specific value. In at least one embodiment, WD is formatted specifically for graphics acceleration module **646** and can be in a form of a graphics acceleration module **646** command, an effective address pointer to a user-defined structure, an effective address pointer to a queue of commands, or any other data structure to describe work to be done by graphics acceleration module **646**.

(184) In at least one embodiment, an AMR value is an AMR state to use for a current process. In at least one embodiment, a value passed to an operating system is similar to an application setting an AMR. In at least one embodiment, if accelerator integration circuit **636** (not shown) and graphics acceleration module **646** implementations do not support a User Authority Mask Override Register (UAMOR), an operating system may apply a current UAMOR value to an AMR value before passing an AMR in a hypervisor call. In at least one embodiment, hypervisor **696** may optionally apply a current Authority Mask Override Register (AMOR) value before placing an AMR into process element **683**. In at least one embodiment, CSRP is one of registers **645** containing an effective address of an area in an application's effective address space **682** for graphics acceleration module **646** to save and restore context state. In at least one embodiment, this pointer is optional if no state is required to be saved between jobs or when a job is preempted. In at least one embodiment, context save/restore area may be pinned system memory.

(185) Upon receiving a system call, operating system **695** may verify that application **680** has registered and been given authority to use graphics acceleration module **646**. In at least one embodiment, operating system **695** then calls hypervisor **696** with information shown in Table 3.

(186) TABLE-US-00003 TABLE 3 OS to Hypervisor Call Parameters

Parameter #	Description
1	A work descriptor (WD)
2	An Authority Mask Register (AMR) value (potentially masked)
3	An effective address (EA)
4	Context Save/Restore Area Pointer (CSRP)
5	A process ID (PID) and optional thread ID (TID)
6	A virtual address (VA) accelerator utilization record pointer (AURP)

Virtual address of storage segment table pointer (SSTP) 7 A logical interrupt service number (LISN)

(187) In at least one embodiment, upon receiving a hypervisor call, hypervisor **696** verifies that operating system **695** has registered and been given authority to use graphics acceleration module **646**. In at least one embodiment, hypervisor **696** then puts process element **683** into a process element linked list for a corresponding graphics acceleration module **646** type. In at least one embodiment, a process element may include information shown in Table 4.

(188) TABLE-US-00004 TABLE 4 Process Element Information

Element #	Description
1	A work descriptor (WD)
2	An Authority Mask Register (AMR) value (potentially masked).
3	An effective address (EA)
Context Save/Restore Area Pointer (CSR)	
4	A process ID (PID) and optional thread ID (TID)
5	A virtual address (VA) accelerator utilization record pointer (AURP)
6	Virtual address of storage segment table pointer (SSTP)
7	A logical interrupt service number (LISN)
8	Interrupt vector table, derived from hypervisor call parameters
9	A state register (SR) value
10	A logical partition ID (LPID)
11	A real address (RA) hypervisor accelerator utilization record pointer
12	Storage Descriptor Register (SDR)

(189) In at least one embodiment, hypervisor initializes a plurality of accelerator integration slice **690** registers **645**.

(190) As illustrated in FIG. **6F**, in at least one embodiment, a unified memory is used, addressable via a common virtual memory address space used to access physical processor memories **601(1)-601(N)** and GPU memories **620(1)-620(N)**. In this implementation, operations executed on GPUs **610(1)-610(N)** utilize a same virtual/effective memory address space to access processor memories **601(1)-601(M)** and vice versa, thereby simplifying programmability. In at least one embodiment, a first portion of a virtual/effective address space is allocated to processor memory **601(1)**, a second portion to second processor memory **601(N)**, a third portion to GPU memory **620(1)**, and so on. In at least one embodiment, an entire virtual/effective memory space (sometimes referred to as an effective address space) is thereby distributed across each of processor memories **601** and GPU memories **620**, allowing any processor or GPU to access any physical memory with a virtual address mapped to that memory.

(191) In at least one embodiment, bias/coherence management circuitry **694A-694E** within one or more of MMUs **639A-639E** ensures cache coherence between caches of one or more host processors (e.g., **605**) and GPUs **610** and implements biasing techniques indicating physical memories in which certain types of data should be stored. In at least one embodiment, while multiple instances of bias/coherence management circuitry **694A-694E** are illustrated in FIG. **6F**, bias/coherence circuitry may be implemented within an MMU of one or more host processors **605** and/or within accelerator integration circuit **636**.

(192) One embodiment allows GPU memories **620** to be mapped as part of system memory, and accessed using shared virtual memory (SVM) technology, but without suffering performance drawbacks associated with full system cache coherence. In at least one embodiment, an ability for GPU memories **620** to be accessed as system memory without onerous cache coherence overhead provides a beneficial operating environment for GPU offload. In at least one embodiment, this arrangement allows software of host processor **605** to setup operands and access computation results, without overhead of tradition I/O DMA data copies. In at least one embodiment, such traditional copies involve driver calls, interrupts and memory mapped I/O (MMIO) accesses that are all inefficient relative to simple memory accesses. In at least one embodiment, an ability to access GPU memories **620** without cache coherence overheads can be critical to execution time of an offloaded computation. In at least one embodiment, in cases with substantial streaming write memory traffic, for example, cache coherence overhead can significantly reduce an effective write bandwidth seen by a GPU **610**. In at least one embodiment, efficiency of operand setup, efficiency of results access, and efficiency of GPU computation may play a role in determining effectiveness of a GPU offload.

(193) In at least one embodiment, selection of GPU bias and host processor bias is driven by a bias tracker data structure. In at least one embodiment, a bias table may be used, for example, which may be a page-granular structure (e.g., controlled at a granularity of a memory page) that includes 1 or 2 bits per GPU-attached memory page. In at least one embodiment, a bias table may be implemented in a stolen memory range of one or more GPU memories **620**, with or without a bias cache in a GPU **610** (e.g., to cache frequently/recently used entries of a bias table). Alternatively, in at least one embodiment, an entire bias table may be maintained within a GPU.

(194) In at least one embodiment, a bias table entry associated with each access to a GPU attached memory **620** is accessed prior to actual access to a GPU memory, causing following operations. In at least one embodiment, local requests from a GPU **610** that find their page in GPU bias are forwarded directly to a corresponding GPU memory **620**. In at least one embodiment, local requests from a GPU that find their page in host bias are forwarded to processor **605** (e.g., over a high-speed link as described herein). In at least one embodiment, requests from processor **605** that find a requested page in host processor bias complete a request like a normal memory read. Alternatively, requests directed to a GPU-biased page may be forwarded to a GPU **610**. In at least one embodiment, a GPU may then transition a page to a host processor bias if it is not currently using a page. In at least one embodiment, a bias state of a page can be changed either by a software-based mechanism, a hardware-assisted software-based mechanism, or, for a limited set of cases, a purely hardware-based mechanism.

(195) In at least one embodiment, one mechanism for changing bias state employs an API call (e.g., OpenCL), which, in turn, calls a GPU's device driver which, in turn, sends a message (or enqueues a command descriptor) to a GPU directing it to change a bias state and, for some transitions, perform a cache flushing operation in a host. In at least one embodiment, a cache flushing operation is used for a transition from host processor **605** bias to GPU bias, but is not for an opposite transition.

(196) In at least one embodiment, cache coherency is maintained by temporarily rendering GPU-biased pages uncacheable by host processor **605**. In at least one embodiment, to access these pages, processor **605** may request access from GPU **610**, which may or may not grant access right away. In at least one embodiment, thus, to reduce communication between processor **605** and GPU **610** it is beneficial to ensure that GPU-biased pages are those which are required by a GPU but not host processor **605** and vice versa.

(197) Hardware structure(s) **115** are used to perform one or more embodiments. FIG. 7A illustrates inference and/or training logic **115** used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **115** are provided below in conjunction with FIGS. 7A and/or 7B.

(198) In at least one embodiment, inference and/or training logic **115** may include, without limitation, code and/or data storage **701** to store forward and/or output weight and/or input/output data, and/or other parameters to configure neurons or layers of a neural network trained and/or used for inferencing in aspects of one or more embodiments. In at least one embodiment, training logic **115** may include, or be coupled to code and/or data storage **701** to store graph code or other software to control timing and/or order, in which weight and/or other parameter information is to be loaded to configure, logic, including integer and/or floating point units (collectively, arithmetic logic units (ALUs)). In at least one embodiment, code, such as graph code, loads weight or other parameter information into processor ALUs based on an architecture of a neural network to which such code corresponds. In at least one embodiment, code and/or data storage **701** stores weight parameters and/or input/output data of each layer of a neural network trained or used in conjunction with one or more embodiments during forward propagation of input/output data and/or weight parameters during training and/or inferencing using aspects of one or more embodiments. In at least one embodiment, any portion of code and/or data storage **701** may be included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory.

(199) In at least one embodiment, any portion of code and/or data storage **701** may be internal or external to one or more processors or other hardware logic devices or circuits. In at least one embodiment, code and/or code and/or data storage **701** may be cache memory, dynamic randomly addressable memory (“DRAM”), static randomly addressable memory (“SRAM”), non-volatile memory (e.g., flash memory), or other storage. In at least one embodiment, a choice of whether code and/or code and/or data storage **701** is internal or external to a processor, for example, or comprising DRAM, SRAM, flash or some other storage type may depend on available storage on-chip versus off-chip, latency requirements of training and/or inferencing functions being performed, batch size of data used in inferencing and/or training of a neural network, or some combination of these factors.

(200) In at least one embodiment, inference and/or training logic **115** may include, without limitation, a code and/or data storage **705** to store backward and/or output weight and/or input/output data corresponding to neurons or layers of a neural network trained and/or used for inferencing in aspects of one or more embodiments. In at least one embodiment, code and/or data storage **705** stores weight parameters and/or input/output data of each layer of a neural network trained or used in conjunction with one or more embodiments during backward propagation of input/output data and/or weight parameters during training and/or inferencing using aspects of one or more embodiments. In at least one embodiment, training logic **115** may include, or be coupled to code and/or data storage **705** to store graph code or other software to control timing and/or order, in which weight and/or other parameter information is to be loaded to configure, logic, including integer and/or floating point units (collectively, arithmetic logic units (ALUs)).

(201) In at least one embodiment, code, such as graph code, causes the loading of weight or other parameter information into processor ALUs based on an architecture of a neural network to which such code corresponds. In at least one embodiment, any portion of code and/or data storage **705** may be included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory. In at least one embodiment, any portion of code and/or data storage **705** may be internal or external to one or more processors or other hardware logic devices or circuits. In at least one embodiment, code and/or data storage **705** may be cache memory, DRAM, SRAM, non-volatile memory (e.g., flash memory), or other storage. In at least one embodiment, a choice of whether code and/or data storage **705** is internal or external to a processor, for example, or comprising DRAM, SRAM, flash memory or some other storage type may depend on available storage on-chip versus off-chip, latency requirements of training and/or inferencing functions being performed, batch size of data used in inferencing and/or training of a neural network, or some combination of these factors.

(202) In at least one embodiment, code and/or data storage **701** and code and/or data storage **705** may be separate storage structures. In at least one embodiment, code and/or data storage **701** and code and/or data storage **705** may be a combined storage structure. In at least one embodiment, code and/or data storage **701** and code and/or data storage **705** may be partially combined and partially separate. In at least one embodiment, any portion of code and/or data storage **701** and code and/or data storage **705** may be included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory.

(203) In at least one embodiment, inference and/or training logic **115** may include, without limitation, one or more arithmetic logic unit(s) (“ALU(s)”) **710**, including integer and/or floating point units, to perform logical and/or mathematical operations based, at least in part on, or indicated by, training and/or inference code (e.g., graph code), a result of which may produce activations (e.g., output values from layers or neurons within a neural network) stored in an activation storage **720** that are functions of input/output and/or weight parameter data stored in code and/or data storage **701** and/or code and/or data storage **705**. In at least one embodiment, activations stored in activation storage **720** are generated according to linear algebraic and or matrix-based mathematics performed by ALU(s) **710** in response to performing instructions or

other code, wherein weight values stored in code and/or data storage **705** and/or data storage **701** are used as operands along with other values, such as bias values, gradient information, momentum values, or other parameters or hyperparameters, any or all of which may be stored in code and/or data storage **705** or code and/or data storage **701** or another storage on or off-chip.

(204) In at least one embodiment, ALU(s) **710** are included within one or more processors or other hardware logic devices or circuits, whereas in another embodiment, ALU(s) **710** may be external to a processor or other hardware logic device or circuit that uses them (e.g., a co-processor). In at least one embodiment, ALUs **710** may be included within a processor's execution units or otherwise within a bank of ALUs accessible by a processor's execution units either within same processor or distributed between different processors of different types (e.g., central processing units, graphics processing units, fixed function units, etc.). In at least one embodiment, code and/or data storage **701**, code and/or data storage **705**, and activation storage **720** may share a processor or other hardware logic device or circuit, whereas in another embodiment, they may be in different processors or other hardware logic devices or circuits, or some combination of same and different processors or other hardware logic devices or circuits. In at least one embodiment, any portion of activation storage **720** may be included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory. Furthermore, inferencing and/or training code may be stored with other code accessible to a processor or other hardware logic or circuit and fetched and/or processed using a processor's fetch, decode, scheduling, execution, retirement and/or other logical circuits.

(205) In at least one embodiment, activation storage **720** may be cache memory, DRAM, SRAM, non-volatile memory (e.g., flash memory), or other storage. In at least one embodiment, activation storage **720** may be completely or partially within or external to one or more processors or other logical circuits. In at least one embodiment, a choice of whether activation storage **720** is internal or external to a processor, for example, or comprising DRAM, SRAM, flash memory or some other storage type may depend on available storage on-chip versus off-chip, latency requirements of training and/or inferencing functions being performed, batch size of data used in inferencing and/or training of a neural network, or some combination of these factors.

(206) In at least one embodiment, inference and/or training logic **115** illustrated in FIG. 7A may be used in conjunction with an application-specific integrated circuit ("ASIC"), such as a TensorFlow® Processing Unit from Google, an inference processing unit (IPU) from Graphcore™, or a Nervana® (e.g., "Lake Crest") processor from Intel Corp. In at least one embodiment, inference and/or training logic **115** illustrated in FIG. 7A may be used in conjunction with central processing unit ("CPU") hardware, graphics processing unit ("GPU") hardware or other hardware, such as field programmable gate arrays ("FPGAs").

(207) FIG. 7B illustrates inference and/or training logic **115**, according to at least one embodiment. In at least one embodiment, inference and/or training logic **115** may include, without limitation, hardware logic in which computational resources are dedicated or otherwise exclusively used in conjunction with weight values or other information corresponding to one or more layers of neurons within a neural network. In at least one embodiment, inference and/or training logic **115** illustrated in FIG. 7B may be used in conjunction with an application-specific integrated circuit (ASIC), such as TensorFlow® Processing Unit from Google, an inference processing unit (IPU) from Graphcore™, or a Nervana® (e.g., "Lake Crest") processor from Intel Corp. In at least one embodiment, inference and/or training logic **115** illustrated in FIG. 7B may be used in conjunction with central processing unit (CPU) hardware, graphics processing unit (GPU) hardware or other hardware, such as field programmable gate arrays (FPGAs). In at least one embodiment, inference and/or training logic **115** includes, without limitation, code and/or data storage **701** and code and/or data storage **705**, which may be used to store code (e.g., graph code), weight values and/or other information, including bias values, gradient information, momentum values, and/or other parameter or hyperparameter information. In at least one embodiment illustrated in FIG. 7B, each of code

and/or data storage **701** and code and/or data storage **705** is associated with a dedicated computational resource, such as computational hardware **702** and computational hardware **706**, respectively. In at least one embodiment, each of computational hardware **702** and computational hardware **706** comprises one or more ALUs that perform mathematical functions, such as linear algebraic functions, only on information stored in code and/or data storage **701** and code and/or data storage **705**, respectively, result of which is stored in activation storage **720**.

(208) In at least one embodiment, each of code and/or data storage **701** and **705** and corresponding computational hardware **702** and **706**, respectively, correspond to different layers of a neural network, such that resulting activation from one storage/computational pair **701/702** of code and/or data storage **701** and computational hardware **702** is provided as an input to a next storage/computational pair **705/706** of code and/or data storage **705** and computational hardware **706**, in order to mirror a conceptual organization of a neural network. In at least one embodiment, each of storage/computational pairs **701/702** and **705/706** may correspond to more than one neural network layer. In at least one embodiment, additional storage/computation pairs (not shown) subsequent to or in parallel with storage/computation pairs **701/702** and **705/706** may be included in inference and/or training logic **115**.

(209) FIG. **8** is a flowchart illustrating an exemplary hardware fault detection process, according to embodiments of the disclosure. As described above, the processor(s) of a feedback control system such as autonomous vehicle **400** may generate multiple instances of its control program such as, e.g., the driving software stack of SoC(s) **404** (Step **800**). These instances are executed in parallel on one or more processors of the feedback control system, such as those described above. System sensors **100** transmit output data that are received by sensor distributor **110** (Step **810**) and transmitted to the various control program instances (Step **820**). As above, sensor data may be transmitted to the various instances in any manner, such as via the various processes of FIG. **2**. Distribution of sensor **100** data may be performed in any manner.

(210) Each instance then generates output actuation commands from its input sensor **100** data. For example, an autonomous vehicle **400** may generate navigation (e.g., steering and engine output) commands calculated from a predetermined path and obstacles detected by one or more of its cameras **468,470,472,498**. These commands are then compared to each other, to determine the presence of a hardware fault (Step **830**). Comparison may be performed in any manner. In some embodiments of the disclosure, hardware faults may be deemed to be present if the outputs of each instance differ from each other by more than some predetermined amount. For example, if target steering outputs of two instances vary by greater than, e.g., 3.0 degrees, a fault may be deemed to have occurred. Any predetermined value or values may be employed for any actuation output commands. As another example, hardware faults may be deemed to be present if any of the last *n* outputs of all instances differ from each other by more than some predetermined amount. If more than two instances are present, comparison may be by any predetermined metric, such as a threshold maximum difference between instance output command values, a threshold variance or deviation among actuation command values, or the like. In some embodiments of the disclosure, comparison may be performed by one or more machine learning models may be trained to detect faults, such as by training to recognize patterns in the control command output of various instances, or patterns in output differences. Any suitable machine learning models may be employed, and these models may be trained in any suitable manner to recognize any desired patterns in control command outputs. For instance, models may be trained to recognize classes of output behavior that correspond to the presence of hardware faults, such as a sudden change in particular outputs, changes in particular outputs, characteristic time-dependent behavior of particular outputs, or the like. Models of embodiments of the disclosure may be trained to recognize any actuator command patterns found to be characteristic of a hardware fault. In some embodiments of the disclosure, comparison of output commands may be by one or more heuristics or rules suitable for identifying such patterns without use of machine learning models. For example, checker **140** may implement a

number of rules that flag the existence of a hardware fault if the output commands satisfy certain criteria, such as differences in output actuator commands of greater than some amount for greater than some specified time period, differences in output actuator commands that grow by more than a specified rate, or the like. Embodiments of the disclosure contemplate any such heuristics or rules. (211) The comparison of Step **830** may determine whether a fault is detected (Step **840**). If no fault is detected, i.e., the output commands of each instance are sufficiently similar to each other, the commands are transmitted to the actuators **150** (Step **860**). Commands may be transmitted in any manner, such as by transmission of commands from one selected instance (with the selected instance being selected in any manner), transmission in a round robin or any other alternating manner, selection of commands from different instances in any manner, averaging of the commands of each instance for a particular actuator or otherwise calculating a command as any function of the output commands of each instance, or the like.

(212) If a hardware fault is detected, the checker **140** may signal the presence of a fault (Step **850**) and/or take any other desired action. As one example, checker **140** may discard actuation commands corresponding to the detected fault, for as long as the fault is determined to exist. As another example, checker **140** may initiate an action of the feedback control system, such as entering the system into a safe or protected mode, disabling particular actuators, maintaining the output of actuators at the last level or setting prior to the hardware fault, or the like. If the feedback control system is an autonomous vehicle **400**, for instance, a checker **140** may initiate an alarm or alert, switch to a different redundant or backup processor if one is available, initiate a vehicle **400** action such as braking, pulling over, shutdown, or the like. Any one or more actions are contemplated.

(213) As used herein, a recitation of “and/or” with respect to two or more elements should be interpreted to mean only one element, or a combination of elements. For example, “element A, element B, and/or element C” may include only element A, only element B, only element C, element A and element B, element A and element C, element B and element C, or elements A, B, and C. In addition, “at least one of element A or element B” may include at least one of element A, at least one of element B, or at least one of element A and at least one of element B. Further, “at least one of element A and element B” may include at least one of element A, at least one of element B, or at least one of element A and at least one of element B.

(214) The subject matter of the present disclosure is described with specificity herein to meet statutory requirements. However, the description itself is not intended to limit the scope of this disclosure. Rather, the inventors have contemplated that the claimed subject matter might also be embodied in other ways, to include different steps or combinations of steps similar to the ones described in this document, in conjunction with other present or future technologies. Moreover, although the terms “step” and/or “block” may be used herein to connote different elements of methods employed, the terms should not be interpreted as implying any particular order among or between various steps herein disclosed unless and except when the order of individual steps is explicitly described.

(215) The foregoing description, for purposes of explanation, used specific nomenclature to provide a thorough understanding of the disclosure. However, it will be apparent to one skilled in the art that the specific details are not required to practice the methods and systems of the disclosure. Thus, the foregoing descriptions of specific embodiments of the present invention are presented for purposes of illustration and description. They are not intended to be exhaustive or to limit the invention to the precise forms disclosed. Many modifications and variations are possible in view of the above teachings. For example, sensor data may be transmitted in any manner or arrangement to various instances of a control program. Any number of control program instances may be present. Fault detection may be by any criteria or manner, and feedback control systems of embodiments of the disclosure may take any actions in response to a detected hardware fault. The embodiments were chosen and described in order to best explain the principles of the invention and

its practical applications, to thereby enable others skilled in the art to best utilize the methods and systems of the disclosure and various embodiments with various modifications as are suited to the particular use contemplated. Additionally, different features of the various embodiments, disclosed or otherwise, can be mixed and matched or otherwise combined so as to create further embodiments contemplated by the disclosure.

Claims

1. A system comprising: processing circuitry to perform operations including: applying, substantially simultaneously, sensor data obtained using one or more sensors to at least two program instances of a same instruction set executed using at least two different respective hardware elements associated with a computing device to generate two or more outputs associated with controlling at least one actuator; determining, using one or more machine learning models (MLMs) and based at least on one or more patterns between the two or more outputs, a result corresponding to hardware fault detection for at least one hardware element of the at least two different respective hardware elements; and performing one or more operations for the computing device based at least on the determining of the result.
2. The system of claim 1, wherein the one or more MLMs include input corresponding to the two or more outputs and provide output indicating the result corresponding to the hardware fault detection.
3. The system of claim 1, wherein the result is based at least on the one or more MLMs detecting the one or more patterns between the two or more outputs.
4. The system of claim 1, wherein the at least two program instances include at least two instances of a control program.
5. The system of claim 1, wherein based at least on the result indicating a first output of the two or more outputs corresponds to a hardware fault, the one or more operations correspond to a second output of the two or more outputs.
6. The system of claim 1, wherein the two or more outputs include a plurality of outputs that correspond to an instance of the at least two program instances.
7. The system of claim 1, wherein the at least two program instances include at least two control stacks of the computing device, and the two or more outputs correspond to control commands determined using the at least two control stacks.
8. The system of claim 1, wherein the applying of the sensor data includes periodically skipping providing units of the sensor data to at least one instance of the at least two program instances to generate the two or more outputs.
9. The system of claim 1, wherein the system is comprised in at least one of: a control system for an autonomous or semi-autonomous machine; a perception system for an autonomous or semi-autonomous machine; a system for performing one or more simulation operations; a system for performing light transport simulation; a system for performing one or more deep learning operations; a system implemented using a robot; a system for presenting at least one of virtual reality content or augmented reality content; a system incorporating one or more virtual machines (VMs); or a system implemented at least partially using cloud computing resources.
10. A method comprising: determining, substantially simultaneously and using at least two program instances of a same instruction set executed using at least two different respective hardware elements associated with a machine, outputs associated with controlling the machine; applying the outputs to one or more machine learning models (MLMs) to generate output data indicating a result corresponding to hardware fault detection for at least one hardware element of the at least two different respective hardware elements based at least on one or more patterns between the outputs; performing one or more operations for the machine based at least on the hardware fault detection.
11. The method of claim 10, wherein the one or more MLMs include input corresponding to the

outputs and provide output indicating the result corresponding to the hardware fault detection.

12. The method of claim 10, wherein the result is based at least on the one or more MLMs detecting the one or more patterns between the outputs.

13. The method of claim 10, wherein the at least two program instances include at least two instances of a control program.

14. The method of claim 10, wherein based at least on the result indicating a first output of the outputs corresponds to a hardware fault, the one or more operations correspond to a second output of the outputs.

15. The method of claim 10, wherein the outputs indicate control commands for the machine and the one or more operations correspond to at least one of the control commands.

16. At least one processor comprising: one or more circuits to: perform hardware fault detection on at least one hardware element of at least two different respective hardware elements of a machine using simulated data generated within a simulation environment, the hardware fault detection performed using one or more machine learning models (MLMs) and based at least on one or more patterns between outputs of at least two program instances of a same instruction set, running on the at least two different respective hardware elements that each process, substantially simultaneously, at least a portion of the simulated data.

17. The at least one processor of claim 16, wherein the one or more MLMs include input corresponding to the outputs and provide output indicating a result corresponding to the hardware fault detection.

18. The at least one processor of claim 16, wherein a result of the hardware fault detection is based at least on the one or more MLMs detecting one or more patterns between the outputs.

19. The at least one processor of claim 16, wherein the simulation environment is rendered using one or more ray-tracing algorithms.

20. The at least one processor of claim 16, wherein the at least one processor is comprised in at least one of: a control system for an autonomous or semi-autonomous machine; a perception system for an autonomous or semi-autonomous machine; a system for performing one or more simulation operations; a system for performing light transport simulation; a system for performing one or more deep learning operations; a system implemented using a robot; a system for presenting at least one of virtual reality content or augmented reality content; a system incorporating one or more virtual machines (VMs); or a system implemented at least partially using cloud computing resources.
