

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12405580
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Townsend-Last; Simon

XML interpreter

Abstract

A system implements an interpreter for extensible markup language (XML) to validate XML produced by large language models (LLMs). The system uses an LLM to generate a first set of XML instructions to perform a first task in an environment communicatively coupled to the system. The system executes the first set of XML instructions and observes a result of the execution. Based on the observed result, the system either uses the LLM to generate a second set of XML instructions to perform a second task, or detects an error in the first set of XML instructions. When an error is detected, the system uses the LLM to modify the first set of XML instructions to correct the error.

Inventors:	Townsend-Last; Simon (New York, NY)
Applicant:	Notion Labs, Inc. (San Francisco, CA)
Family ID:	96175055
Assignee:	Notion Labs, Inc. (San Francisco, CA)
Appl. No.:	18/655037
Filed:	May 03, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20250217227 A1	Jul. 03, 2025

Related U.S. Application Data

us-provisional-application US 63616450 20231229

Publication Classification

Int. Cl.: G06F11/00 (20060101); G06F11/07 (20060101); G06F16/955 (20190101)

U.S. Cl.:

CPC G06F11/0793 (20130101); G06F11/0736 (20130101); G06F16/955 (20190101);

Field of Classification Search

CPC: G06F (11/0793); G06F (11/0736); G06F (16/955)

USPC: 714/1-57

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7512840	12/2008	Martin	717/124	G06F 40/226
12118320	12/2023	Gelfenbeyn	N/A	G06F 40/30
2009/0055686	12/2008	Tsang	714/33	G06F 11/3698
2019/0332680	12/2018	Wang	N/A	G10L 15/07
2022/0101839	12/2021	George et al.	N/A	N/A
2023/0125194	12/2022	Mertens et al.	N/A	N/A
2025/0045492	12/2024	Koiek Akino	N/A	G06F 30/27
2025/0111220	12/2024	Jones	N/A	G06N 3/08

OTHER PUBLICATIONS

ISA, International Search Report for International Application No. PCT/US2024/036450, mailed on Sep. 23, 2024, 4 pages. cited by applicant

Primary Examiner: Butler; Sarai E

Attorney, Agent or Firm: Perkins Coie LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION (1) This application claims the benefit of U.S. Provisional Patent Application No. 63/616,450, filed Dec. 29, 2023, which is incorporated herein by reference in its entirety.

BACKGROUND

(1) Many industries are turning to artificial intelligence tools to automate tasks that previously required significant human labor or were infeasible or impossible for humans to perform. However, despite advancement of these tools, integrating them into some types of environments has proven challenging. Existing tools, for example, lack the inherent capacity to autonomously comprehend and navigate structured software environments without extensive manual guidance. These limitations hamper the ability of artificial intelligence tools to perform tasks seamlessly and efficiently within these environments.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) Reference will now be made, by way of example, to the accompanying drawings which show example embodiments of the present application, and in which:
- (2) FIG. 1 is a high-level block diagram illustrating an environment in which an artificial intelligence assistant operates, according to some implementations.
- (3) FIGS. 2A-2E illustrate example interfaces for a user to interact with an artificial intelligence assistant.
- (4) FIG. 3 is a block diagram illustrating a platform, which may be used to implement examples of the present disclosure.
- (5) FIG. 4 is a block diagram illustrating a hierarchical organization of pages in a workspace.
- (6) FIG. 5 is a block diagram of a transformer neural network, which may be used in examples of the present disclosure.
- (7) FIG. 6 is a flowchart illustrating a process for controlling an environment using an artificial intelligence assistant, according to some implementations.
- (8) FIG. 7 is a flowchart illustrating a process for implementing an XML interpreter to validate instructions that are written by or with the assistance of a large language model, according to some implementations.
- (9) FIG. 8 is a block diagram that illustrates an example of a computer system in which at least some operations described herein can be implemented.
- (10) The technologies described herein will become more apparent to those skilled in the art by studying the Detailed Description in conjunction with the drawings. Embodiments or implementations describing aspects of the invention are illustrated by way of example, and the same references can indicate similar elements. While the drawings depict various implementations for the purpose of illustration, those skilled in the art will recognize that alternative implementations can be employed without departing from the principles of the present technologies. Accordingly, while specific implementations are shown in the drawings, the technology is amenable to various modifications.

DETAILED DESCRIPTION

- (11) Artificial intelligence tools provide many beneficial features, from automating routine tasks to performing complex analyses of large datasets. However, integrating these tools into an environment can prove challenging because the tools cannot autonomously navigate these environments. The present technology provides a structured workflow by which an artificial intelligence tool can interact with an environment to perform tasks in the environment. In this workflow, a computer system implements an artificial intelligence tool that performs tasks in the environment in response to natural language instructions received from a user. The tool leverages a large language model to generate computer-readable instruction sets that can perform tasks autonomously in any environment.
- (12) According to implementations herein, a system receives a natural language instruction to perform a task within an environment that is communicatively coupled to the system. Based on the received instruction, the system generates a computer-readable input that includes a context of the environment and a computer-readable form of the received instruction. The system sends the computer-readable input to a large language model (LLM) to cause the LLM to generate a set of computer program code to perform the task. The computer program code is executed to perform the task in the environment.
- (13) Some aspects disclosed herein relate to a system that validates code output by an LLM in order to perform tasks in an environment. The system can use an LLM to generate a first set of extensible markup language (XML) instructions to perform a first task in an environment

communicatively coupled to the system. The system executes the first set of XML instructions and observes a result of the execution. Based on the observed result of the execution, the system either uses the LLM to generate a second set of XML instructions to perform a second task in the environment, or detects an error in the first set of XML instructions. When an error is detected, the system uses the LLM to modify the first set of XML instructions to correct the error.

(14) The description and associated drawings are illustrative examples and are not to be construed as limiting. This disclosure provides certain details for a thorough understanding and enabling description of these examples. One skilled in the relevant technology will understand, however, that the invention can be practiced without many of these details. Likewise, one skilled in the relevant technology will understand that the invention can include well-known structures or features that are not shown or described in detail, to avoid unnecessarily obscuring the descriptions of examples.

(15) Artificial Intelligence Assistant Environment

(16) FIG. 1 is a high-level block diagram illustrating an environment **100** in which an artificial intelligence assistant operates, according to some implementations. As shown in FIG. 1, the environment **100** can include a controlled environment **110**, an assistant **120**, and a large language model (LLM) **130**.

(17) The controlled environment **110** is a physical or virtual environment that is operated by one or more computing systems. In an example, the controlled environment **110** is a virtual environment that is accessed via user devices, such as a website, a web application, or a native application. Other example controlled environments **110** include physical systems that are operated by controllers coupled to computing systems, such as manufacturing or testing facilities that employ robotic systems to perform tasks. An example controlled environment **110**, in the form of a data and project management platform, is described with respect to FIG. 3.

(18) A user can interact with the controlled environment **110** via a user device. When the controlled environment **110** is a virtual environment, for example, the user device can access and display pages or content from the virtual environment to a user. A user can read or edit the environment's content from the user device. Users can also interact with the assistant **120** via the user devices to automate tasks in the environment **110**.

(19) The LLM **130** includes one or more language models that are configured to generate text-based outputs in response to prompts. The LLM **130** can include any commercially available or custom models, or a set or ensemble of two or more models. Example features of LLMs are described with respect to FIG. 5.

(20) The LLM **130** can be trained to manipulate computer-readable instructions and operate application programming interfaces (APIs) in order to perform tasks in the controlled environment **110**. During training, the LLM **130** can be provided with example transcripts that include sample instructions from a user and corresponding code to implement a task in the controlled environment. Training the LLM **130** can include, for example, preference model training that trains the LLM to predict a human preference based on ranked or contrastive pairs, supervised learning that uses human feedback data or synthetic data to train the LLM to predict a next token in a sequence, or reinforcement learning that penalizes the LLM for outputs that do not satisfy specified preference requirements.

(21) The assistant **120** is a computer system or software application that communicates with the controlled environment **110** and a user to perform tasks in the environment **110** based on natural language inputs received from the user. In various implementations, the assistant **120** can communicate with the computer systems that implement or control the environment **110** via a network such as a local network or the Internet, or can be integrated into a device or system that controls the environment. The assistant **120** can interface with the controlled environment **110** to perform observations of the environment and to effect actions within the environment.

(22) An example interface for a user to interact with the assistant **120** is illustrated in FIGS. 2A-2E.

The example in FIG. 2A illustrates a page 205 within a data management platform, as an example controlled environment 110. The page 205 includes a text block 210 entitled “Shopping list” with bullet points under the title.

(23) A user of the data management platform can interact with the assistant 120 to perform tasks associated with the page 205. Any of a variety of mechanisms for accessing the functionality of the assistant 120 can be provided within, or associated with, the environment 110. In the example in FIG. 2A, a user can invoke a text box 215 to input a natural language instruction. For example, the text box 215 can be displayed within the page 205 or in a modal window or sidebar associated with the page 205.

(24) In one example, illustrated in FIG. 2B, a user interacts with the assistant 120 via a chat-like interface 220. The example of FIG. 2B illustrates that the user has queried the assistant for information about the page, asking “what is the first item on the shopping list?” In response, the assistant 120 outputs the text, “The first item on your shopping list is Apples.” The response from the assistant 120 can be provided within the chat interface, in some implementations.

(25) A user can also interact with the assistant 120 within the context of the page 205 to modify the page's content. In FIG. 2C, for example, a text box 215 is provided as a new block on the page 205. In FIG. 2D, the user has entered the natural language instruction “Below this block, insert a recipe using at least three items from the shopping list.” In response, the assistant 120 generates a recipe and inserts the generated text into a new block 230 on the page, as shown in FIG. 2E, replacing the user's natural language instruction.

(26) Users can interact with the assistant 120 in ways other than those illustrated in FIGS. 2A-2E. For example, a user can chat with an assistant 120 via an application that is separate from the controlled environment 110, including third-party social media applications or chat, instant messaging, or collaboration platforms such as Slack, Webex, or Microsoft Teams. Likewise, a user can interact with an assistant 120 within a first application (e.g., a Slack thread) to perform tasks in a second application linked to the first application (e.g., a data management platform linked to the Slack thread). Alternatively, a custom application that is accessible to users via a user's computing device can integrate with the controlled environment 110 to effect changes in the environment based on user inputs at the custom application.

(27) As will be described further below, the assistant 120 generates computer program code in order to perform observations of the controlled environment 110 and to effects actions within the environment. During a user's interactions with the assistant 120, the assistant 120 generates a transcript that represents a sequence of inputs, outputs, or observations and that maintains a persistent state of this sequence. The transcript includes computer-readable inputs and computer program code that is executed by the assistant 120 to perform tasks related to the controlled environment 110. For example, the transcript includes a set of extensible markup language (XML), JavaScript, or a combination of XML and JavaScript. Provided below is a portion of an example transcript generated during the chat interaction illustrated in FIG. 2B, in which the assistant observes and outputs an identification of the first item in a list entitled “Shopping List:”

(28) TABLE-US-00001 [{ “id”: 0, “type”: “context”, “context”: (...) }, { “id”: 1, “type”: “assistant”, “value”: “<load-page id=\\“0\\”/>” }, { “type”: “observation”, “observationType”: “page”, “pageId”: “0”, “value”: “<page id=\\“0\\”><property-title name=\\“Title\\”/><text id=\\“2\\”>Shopping list:</text><uli id=\\“3\\”>Apples</uli><uli id=\\“4\\”>” }, { “id”: 2, “type”: “human”, “value”: “<chat><text>What's the first item on my shopping list?</text></chat>” }]

(29) The transcript can include a series of steps, where each step includes computer program instructions associated with user steps or assistant steps. User steps in the transcript can include computer-readable inputs that are generated based on natural language instructions received from a user. The computer-readable inputs can include translating a user's natural language input into a

computer-readable form. For example, a user enters, at a text entry box associated with the assistant **120**, the natural language instruction, “Please update this page to include this information: {information}.” When generating the transcript, the assistant **120** translates this instruction into the following XML: <chat><text>Please update this page to incorporate this information: {information}</text></chat>

(30) At least some user steps can also include a context of the controlled environment **110** at the time a user input was received. The context can include, for example, a state of the environment **110** (e.g., a page the user is viewing or a thread in a collaboration application with which a user is interacting), a date and time of the user input, previous interactions between the user and the assistant **120**, or other information that enables the assistant **120** to perform tasks. In some implementations, context is included in any user step that initiates a new interaction with the assistant **120**. Context may not be included in user steps of the transcript that continue a prior interaction with the assistant **120**, such as the user continuing a chat conversation with the assistant or the user asking the assistant to revise content that the assistant previously generated.

(31) Assistant steps in the transcript include observation steps and action steps. Observations of the controlled environment **110** can include computer program code that, when executed, causes the assistant to observe a value or state of the controlled environment **110**. For example, an observation step can return the name of an element on a page or a block of text that matches a query, retrieve a value from a table, identify a relevant document in a document repository, observe properties of elements on a page, determine a last edit date of a document, and so forth.

(32) Action steps in the transcript can include computer program code that, when executed, modifies properties or content within the controlled environment **110**. For example, when the controlled environment **110** is a virtual environment such as a data management platform, action steps can include commands to load a page, insert content before or after a specified point on a page, insert content inside of another content block, move content on a page, delete content, or set or modify properties or attributes of items on a page.

(33) Some tasks requested by users are deterministic tasks for which the task result is expected to be a certain, predictable output. For example, reading a value from the environment **110**, writing a certain value to the environment, or performing a mathematical operation are deterministic tasks. Other tasks are non-deterministic tasks, such as if a user requests a summary of a document or a recipe that includes items from a shopping list. To perform such deterministic or non-deterministic tasks, at least some implementations of the assistant **120** can cause the LLM to generate XML or some other code, like JavaScript.

(34) In an example, a user inputs a natural language instruction to add content to a structured digital environment, such as a data management platform, where the instruction specifies a location within the environment at which the content should be added. The assistant **120** processes the user's instruction and adds a computer-readable input for the instruction to a transcript. The assistant **120** then generates a prompt to the LLM **130** to generate computer-readable code to perform the requested task. In response, the LLM **130** generates code that is configured to use context of the environment, provided with the computer-readable input or referenced earlier in the transcript, to identify the location within the structured digital environment at which the content is to be inserted. The code generated by the LLM **130** also is executable to cause the content to be written to the structured digital environment at the identified location. The code generated by the LLM **130** is added to the transcript and executed by the assistant **120** to complete the task. In another example, the task requested to be performed by the assistant **120** can be either deterministic or non-deterministic.

(35) In another example, a user inputs a natural language instruction at a data management platform that instructs the assistant **120** to generate content for output via a chat thread (e.g., on a third-party messaging or collaboration platform). To write content to a chat thread, the assistant **120** can collaborate with the LLM **130** on a transcript in a manner like that described in the example above.

However, when generating a computer-readable input based on the natural language instruction, the assistant **120** can include a context of the chat thread in the input to enable the assistant to write content to the specified chat thread based on the code generated by the LLM **130**. Similarly, if a user inputs a natural language instruction within a chat thread that instructs the assistant **120** to perform a task in a data management platform, the computer-readable input generated by the assistant **120** can include a context of the data management platform.

(36) The assistant **120** can also validate outputs from the LLM **130** to ensure that tasks are performed correctly. The LLM **130** may at times output incorrect code, for example by hallucinating APIs or libraries that do not exist in the language in which the code is written or that are not accessible to the assistant **120**. The LLM **130** may also employ improper syntax, generate incorrect data types, fail to fully implement algorithms, or otherwise generate code with bugs, logic errors, or other problems. Errors in the code can also arise based on incorrect inputs by a human user or based on changes to the controlled environment **110**. For example, an error may arise when a user requests an update to a table entitled “Quarterly Finances” but the controlled environment **110** has two tables with the same name.

(37) To validate the computer program code output by the LLM **130**, the assistant **120** executes instructions and observes results of these executions. Generally, if an observed result matches an expected task response, the assistant **120** determines that the instructions output by the LLM **130** are valid. If the assistant **120** detects an error in an observed result, the assistant **120** causes the LLM **130** to produce new code to correct the error before a task response is finalized. A process for the assistant **120** to validate code written by the LLM **130** is described further with respect to FIG. 7.

(38) Block Data Model

(39) Some implementations of the controlled environment **110** are structured according to a block data model (“block model”). According to these implementations, the blocks are dynamic units of information that can be transformed into other block types and move across workspaces in response to either user inputs or based on automated tasks performed by a computing system (such as the assistant **120**). The block model allows users or the computer system to customize how information is moved, organized, and shared. Hence, blocks contain information but are not siloed.

(40) Blocks are singular pieces that represent all units of information inside an editor. In one example, text, images, lists, a row in a database, etc., are all blocks in a workspace. The attributes of a block determine how that information is rendered and organized. Every block can have attributes including an identifier (ID), properties, and type. Each block is uniquely identifiable by its ID. The properties can include a data structure containing custom attributes about a specific block. An example of a property is “title,” which stores text content of block types such as paragraphs, lists, and the title of a page. More elaborate block types require additional or different properties, such as a page block in a database with user-defined properties. Every block can have a type, which defines how a block is displayed and how the block's properties are interpreted.

(41) A block has attributes that define its relationship with other blocks. For example, the attribute “content” is an array (or ordered set) of block IDs representing the content inside a block, such as nested bullet items in a bulleted list or the text inside a toggle. The attribute “parent” is the block ID of a block's parent, which can be used for permissions. Blocks can be combined with other blocks to track progress and hold all project information in one place.

(42) A block type is what specifies how the block is rendered in a user interface (UI), and the block's properties and content are interpreted differently depending on that type. Changing the type of a block does not change the block's properties or content—it only changes the type attribute. The information is thus rendered differently or even ignored if the property is not used by that block type. Decoupling property storage from block type allows for efficient transformation and changes to rendering logic and is useful for collaboration.

(43) Blocks can be nested inside of other blocks (e.g., infinitely nested sub-pages inside of pages).

The content attribute of a block stores the array of block IDs (or pointers) referencing those nested blocks. Each block defines the position and order in which its content blocks are rendered. This hierarchical relationship between blocks and their render children are referred to herein as a “render tree.” In one example, page blocks display their content in a new page, instead of rendering it indented in the current page. To see this content, a user would need to click into the new page.

(44) In the block model, indentation is structural (e.g., reflects the structure of the render tree). In other words, when a user indents something, the user is manipulating relationships between blocks and their content, not just adding a style. For example, pressing Indent in a content block can add that block to the content of the nearest sibling block in the content tree.

(45) Blocks can inherit permissions of blocks in which they are located (which are above them in the tree). Consider a page: to read its contents, a user must be able to read the blocks within that page. However, there are two reasons one cannot use the content array to build the permissions system. First, blocks are allowed to be referenced by multiple content arrays to simplify collaboration and a concurrency model. But because a block can be referenced in multiple places, it is ambiguous which block it would inherit permissions from. The second reason is mechanical. To implement permission checks for a block, one needs to look up the tree, getting that block's ancestors all the way up to the root of the tree (which is the workspace). Trying to find this ancestor path by searching through all blocks' content arrays is inefficient, especially on the client. Instead, the model uses an “upward pointer”—the parent attribute—for the permission system. The upward parent pointers and the downward content pointers mirror each other.

(46) A block's life starts on the client. When a user takes an action in the interface—typing in the editor, dragging blocks around a page—these changes are expressed as operations that create or update a single record. The “records” refer to persisted data, such as blocks, users, workspaces, etc. Because many actions usually change more than one record, operations are batched into transactions that are committed (or rejected) by the server as a group.

(47) Creating and updating blocks can be performed by, for example, pressing Enter on a keyboard. First, the client defines all the initial attributes of the block, generating a new unique ID, setting the appropriate block type (to_do), and filling in the block's properties (an empty title, and checked. The client builds operations to represent the creation of a new block with those attributes. New blocks are not created in isolation: blocks are also added to their parent's content array, so they are in the correct position in the content tree. As such, the client also generates an operation to do so. All these individual change operations are grouped into a transaction. Then, the client applies the operations in the transaction to its local state. New block objects are created in memory and existing blocks are modified. In native apps, the model caches all records that are accessed locally in an LRU (least recently used) cache on top of SQLite or IndexedDB, referred to as RecordCache. When records are changed on a native app, the model also updates the local copies in RecordCache. The editor re-renders to draw the newly created block onto the display. At the same time, the transaction is saved into TransactionQueue, the part of the client responsible for sending all transactions to the model's servers so that the data is persisted and shared with collaborators. TransactionQueue stores transactions safely in IndexedDB or SQLite (depending on the platform) until they are persisted by the server or rejected.

(48) A block can be saved on a server to be shared with others. Usually, TransactionQueue sits empty, so the transaction to create the block is sent to the server in an application programming interface (API) request. In one example, the transaction data is serialized to JSON and posted to the/saveTransactions API endpoint.

(49) Save Transactions gets the data into source-of-truth databases, which store all block data as well as other kinds of persisted records. Once the request reaches the API server, all the blocks and parents involved in the transaction are loaded. This gives a “before” picture in memory. The block model duplicates the “before” data that had just been loaded in memory. Next, the block model applies the operations in the transaction to the new copy to create the “after” data. Then the model

uses both “before” and “after” data to validate the changes for permissions and data coherency. If everything checks out, all created or changed records are committed to the database-meaning the block has now officially been created. At this point, a “success” HTTP response to the original API request is sent by the client. This confirms that the client knows the transaction was saved successfully and that it can move on to saving the next transaction in the TransactionQueue. In the background, the block model schedules additional work depending on the kind of change made for the transaction. For example, the block model can schedule version history snapshots and indexing block text for a Quick Find function. The block model also notifies MessageStore, which is a real-time updates service, about the changes that were made.

(50) The block model provides real-time updates to, for example, almost instantaneously show new blocks to members of a teamspace. Every client can have a long-lived WebSocket connection to the MessageStore. When the client renders a block (or page, or any other kind of record), the client subscribes to changes of that record from MessageStore using the WebSocket connection. When a team member opens the same page, the member is subscribed to changes of all those blocks. After changes have been made through the saveTransactions process, the API notifies MessageStore of new recorded versions. MessageStore finds client connections subscribed to those changing records and passes on the new version through their WebSocket connection. When a team member's client receives version update notifications from MessageStore, it verifies that version of the block in its local cache. Because the versions from the notification and the local block are different, the client sends a syncRecordValues API request to the server with the list of outdated client records. The server responds with the new record data. The client uses this response data to update the local cache with the new version of the records, then re-renders the user interface to display the latest block data.

(51) Blocks can be shared instantaneously with collaborators. In one example, a page is loaded using only local data. On the web, block data is pulled from being in memory. On native apps, loading blocks that are not in memory are loaded from the RecordCache persisted storage. However, if missing block data is needed, the data is requested from an API. The API method for loading the data for a page is referred to herein as loadPageChunk; it descends from a starting point (likely the block ID of a page block) down the content tree and returns the blocks in the content tree plus any dependent records needed to properly render those blocks. Several layers of caching for loadPageChunk are used, but in the worst case, this API might need to make multiple trips to the database as it recursively crawls down the tree to find blocks and their record dependencies. All data loaded by loadPageChunk is put into memory (and saved in the RecordCache if using the app). Once the data is in memory, the page is laid out and rendered using React.

(52) Software Platform

(53) FIG. 3 is a block diagram of a platform 300, aspects of which can function as an example controlled environment 110. The platform 300 provides users with an all-in-one workspace for data and project management. The platform 300 can include a user application 302, an AI tool 304, and a server 306. The user application 302, the AI tool 304, and the server 306 are in communication with each other via a network.

(54) In some implementations, the user application 302 is a cross-platform software application configured to work on several computing platforms and web browsers. The user application 302 can include a variety of templates. A template refers to a prebuilt page that a user can add to a workspace within the user application 302. The templates can be directed to a variety of functions. Exemplary templates include a docs template 308, a wikis template 310, a projects template 312, and a meeting and calendar template 314. In some implementations, a user can generate, save, and share customized templates with other users.

(55) The user application 302 templates can be based on content “blocks.” For example, the templates of the user application 302 include a predefined and/or pre-organized set of blocks that can be customized by the user. Blocks are content containers within a template that can include

text, images, objects, tables, maps, and/or other pages (e.g., nested pages or sub-pages). Blocks can be assigned to certain properties. The blocks are defined by boundaries having dimensions. The boundaries can be visible or non-visible for users. For example, a block can be assigned as a text block (e.g., a block including text content), a heading block (e.g., a block including a heading) or a sub-heading block having a specific location and style to assist in organizing a page. A block can be assigned as a list block to include content in a list format. A block can be assigned as an AI prompt block (also referred to as a “prompt block”) that enables a user to provide instructions (e.g., prompts) to the AI tool **304** to perform functions. A block can also be assigned to include audio, video, or image content.

(56) A user can add, edit, and remove content from the blocks. The user can also organize the content within a page by moving the blocks around. In some implementations, the blocks are shared (e.g., by copying and pasting) between the different templates within a workspace. For example, a block embedded within multiple templates can be configured to show edits synchronously.

(57) The docs template **308** is a document generation and organization tool that can be used for generating a variety of documents. For example, the docs template **308** can be used to generate pages that are easy to organize, navigate, and format. The wikis template **310** is a knowledge management application having features similar to the pages generated by the docs template **308** but that can additionally be used as a database. The wikis template **310** can include, for example, tags configured to categorize pages by topic and/or include an indication of whether the provided information is verified to indicate its accuracy and reliability. The projects template **312** is a project management and note-taking software tool. The projects template **312** can allow the users, either as individuals or as teams, to plan, manage, and execute projects in a single forum. The meeting and calendar template **314** is a tool for managing tasks and timelines. In addition to traditional calendar features, the meeting and calendar template **314** can include blocks for categorizing and prioritizing scheduled tasks, generating to-do and action item lists, tracking productivity, etc. The various templates of the user application **302** can be included under a single workspace and include synchronized blocks. For example, a user can update a project deadline on the projects template **312**, which can be automatically synchronized to the meeting and calendar template **314**. The various templates of the user application **302** can be shared within a team, allowing multiple users to modify and update the workspace concurrently.

(58) The AI tool **304** is an integrated AI assistant that enables AI-based functions for the user application **302**. In one example, the AI tool **304** is based on a neural network architecture, such as the transformer **512** described in FIG. 5. The AI tool **304** can interact with blocks embedded within the templates on a workspace of the user application **302**. For example, the AI tool **304** can include a writing assistant tool **316**, a knowledge management tool **318**, a project management tool **320**, and a meeting and scheduling tool **322**. The different tools of the AI tool **304** can be interconnected and interact with different blocks and templates of the user application **302**.

(59) The writing assistant tool **316** can operate as a generative AI tool for creating content for the blocks in accordance with instructions received from a user. Creating the content can include, for example, summarizing, generating new text, or brainstorming ideas. For example, in response to a prompt received as a user input that instructs the AI to describe what the climate is like in New York, the writing assistant tool **316** can generate a block including a text that describes the climate in New York. As another example, in response to a prompt that requests ideas on how to name a pet, the writing assistant tool **316** can generate a block including a list of creative pet names. The writing assistant tool **316** can also operate to modify existing text. For example, the writing assistant can shorten, lengthen, or translate existing text, correct grammar and typographical errors, or modify the style of the text (e.g., a social media style versus a formal style).

(60) The knowledge management tool **318** can use AI to categorize, organize, and share knowledge included in the workspace. In some implementations, the knowledge management tool **318** can

operate as a question-and-answer assistant. For example, a user can provide instructions on a prompt block to ask a question. In response to receiving the question, the knowledge management tool **318** can provide an answer to the question, for example, based on information included in the wikis template **310**. The project management tool **320** can provide AI support for the projects template **312**. The AI support can include auto filling information based on changes within the workspace or automatically track project development. For example, the project management tool **320** can use AI for task automation, data analysis, real-time monitoring of project development, allocation of resources, and/or risk mitigation. The meeting and scheduling tool **322** can use AI to organize meeting notes, unify meeting records, list key information from meeting minutes, and/or connect meeting notes with deliverable deadlines.

(61) The server **306** can include various units (e.g., including compute and storage units) that enable the operations of the AI tool **304** and workspaces of the user application **302**. The server **306** can include an integrations unit **324**, an application programming interface (API) **328**, databases **326**, and an administration (admin) unit **330**. The databases **326** are configured to store data associated with the blocks. The data associated with the blocks can include information about the content included in the blocks, the function associated with the blocks, and/or any other information related to the blocks. The API **328** can be configured to communicate the block data between the user application **302**, the AI tool **304**, and the databases **326**. The API **328** can also be configured to communicate with remote server systems, such as AI systems. For example, when a user performs a transaction within a block of a template of the user application **302** (e.g., in a docs template **308**), the API **328** processes the transaction and saves the changes associated with the transaction to the database **326**. The integrations unit **324** is a tool connecting the platform **200** with external systems and software platforms. Such external systems and platforms can include other databases (e.g., cloud storage spaces), messaging software applications, or audio or video conference applications. The administration unit **330** is configured to manage and maintain the operations and tasks of the server **306**. For example, the administration unit **330** can manage user accounts, data storage, security, performance monitoring, etc.

(62) FIG. **4** is a block diagram illustrating a hierarchical organization of pages in a workspace. As described with respect to the block data model of the present technology, a workspace can include multiple pages (e.g., page blocks). The pages (e.g., including parent pages and child or nested pages) can be arranged hierarchically within the workspace or one or more teamspace, as shown in FIG. **4**. The page can include a block such as tabs, lists, images, tables, etc.

(63) A teamspace can refer to a collaborative space associated with a team or an organization that is hierarchically below a workspace. For example, a workspace can include a teamspace accessible by all users of an organization and multiple teamspace that are accessible by users of different teams. Accessibility generally refers to creating, editing, and/or viewing content (e.g., pages) included in the workspace or the one or more teamspace.

(64) In the hierarchical organization illustrated in FIG. **4**, a parent page (e.g., “Parent Page”) is located hierarchically below the workspace or a teamspace. The parent page includes three children pages (e.g., “Page 1,” “Page 2,” and “Page 3”). Each of the child pages can further include subpages (e.g., “Page 2 Child” which is a grandchild of “Parent Page” and child of “Page 2”). The “Content” arrows in FIG. **4** indicate the relationship between the parents and children while the “Parent” arrows indicate the inheritance of access permissions. The child pages inherit access permission from the (immediate) parent page under which they are located hierarchically (e.g., which is above them in the tree). For example, “Page 2” inherited the access permission of the “Parent page” as a default when it was created under its parent page. Similarly, “Page 2 Child” inherited the access permission of the parent page as a default when it was created under its parent page. “Parent Page,” “Page 2,” and “Page 2 Child” thereby have the same access permission within the workspace.

(65) The relationships and organization of the content can be modified by changing the location of

the pages. For example, when a child page is moved to be under a different parent, the child page's access permission modifies to correspond to the access permission of the new parent. Also, when the access permission of "Parent Page" is modified, the access permission of "Page 1," "Page 2," and "Page 3" can be automatically modified to correspond to the access permission of "Parent Page" based on the inheritance character of access permissions.

(66) In contrast, however, a user can modify the access permission of the children independently of their parents. For example, the user can modify the access permission of "Page 2 Child" in FIG. 4 so that it is different from the access permission of "Page 2" and "Parent Page." The access permission of "Page 2 Child" can be modified to be broader or narrower than the access permission of its parents. As an example, "Page 2 Child" can be shared on the internet while "Page 2" is only shared internally to the users associated with the workspace. As another example, "Page 2 Child" can be shared only with an individual user while "Page 2" is shared with a group of users (e.g., a team of the organization associated with the workspace). In some implementations, the hierarchical inheritance of the access permissions described herein can be modified from the previous description. For example, the access permissions of all the pages (parent and children) can be defined as independently changeable.

(67) Transformer for Neural Network

(68) To assist in understanding the present disclosure, some concepts relevant to neural networks and machine learning (ML) are discussed herein. Generally, a neural network comprises a number of computation units (sometimes referred to as "neurons"). Each neuron receives an input value and applies a function to the input to generate an output value. The function typically includes a parameter (also referred to as a "weight") whose value is learned through the process of training. A plurality of neurons may be organized into a neural network layer (or simply "layer") and there may be multiple such layers in a neural network. The output of one layer may be provided as input to a subsequent layer. Thus, input to a neural network may be processed through a succession of layers until an output of the neural network is generated by a final layer. This is a simplistic discussion of neural networks and there may be more complex neural network designs that include feedback connections, skip connections, and/or other such possible connections between neurons and/or layers, which are not discussed in detail here.

(69) A deep neural network (DNN) is a type of neural network having multiple layers and/or a large number of neurons. The term DNN can encompass any neural network having multiple layers, including convolutional neural networks (CNNs), recurrent neural networks (RNNs), multilayer perceptrons (MLPs), Generative Adversarial Networks (GANs), Variational Autoencoders (VAEs), and Auto-regressive Models, among others.

(70) DNNs are often used as ML-based models for modeling complex behaviors (e.g., human language, image recognition, object classification, etc.) in order to improve the accuracy of outputs (e.g., more accurate predictions) such as, for example, as compared with models with fewer layers. In the present disclosure, the term "ML-based model" or more simply "ML model" may be understood to refer to a DNN. Training an ML model refers to a process of learning the values of the parameters (or weights) of the neurons in the layers such that the ML model is able to model the target behavior to a desired degree of accuracy. Training typically requires the use of a training dataset, which is a set of data that is relevant to the target behavior of the ML model.

(71) As an example, to train an ML model that is intended to model human language (also referred to as a "language model"), the training dataset may be a collection of text documents, referred to as a "text corpus" (or simply referred to as a "corpus"). The corpus may represent a language domain (e.g., a single language), a subject domain (e.g., scientific papers), and/or may encompass another domain or domains, be they larger or smaller than a single language or subject domain. For example, a relatively large, multilingual, and non-subject-specific corpus can be created by extracting text from online webpages and/or publicly available social media posts. Training data can be annotated with ground truth labels (e.g., each data entry in the training dataset can be paired

with a label) or may be unlabeled.

(72) Training an ML model generally involves inputting into an ML model (e.g., an untrained ML model) training data to be processed by the ML model, processing the training data using the ML model, collecting the output generated by the ML model (e.g., based on the inputted training data), and comparing the output to a desired set of target values. If the training data is labeled, the desired target values may be, e.g., the ground truth labels of the training data. If the training data is unlabeled, the desired target value may be a reconstructed (or otherwise processed) version of the corresponding ML model input (e.g., in the case of an autoencoder), or can be a measure of some target observable effect on the environment (e.g., in the case of a reinforcement learning agent). The parameters of the ML model are updated based on a difference between the generated output value and the desired target value. For example, if the value outputted by the ML model is excessively high, the parameters may be adjusted so as to lower the output value in future training iterations. An objective function is a way to quantitatively represent how close the output value is to the target value. An objective function represents a quantity (or one or more quantities) to be optimized (e.g., minimize a loss or maximize a reward) in order to bring the output value as close to the target value as possible. The goal of training the ML model typically is to minimize a loss function or maximize a reward function.

(73) The training data can be a subset of a larger data set. For example, a data set may be split into three mutually exclusive subsets: a training set, a validation (or cross-validation) set, and a testing set. The three subsets of data may be used sequentially during ML model training. For example, the training set may be first used to train one or more ML models, each ML model, e.g., having a particular architecture, having a particular training procedure, being describable by a set of model hyperparameters, and/or otherwise being varied from the other of the one or more ML models. The validation (or cross-validation) set may then be used as input data into the trained ML models to, e.g., measure the performance of the trained ML models and/or compare performance between them. Where hyperparameters are used, a new set of hyperparameters can be determined based on the measured performance of one or more of the trained ML models, and the first step of training (e.g., with the training set) may begin again on a different ML model described by the new set of determined hyperparameters. In this way, these steps can be repeated to produce a more performant trained ML model. Once such a trained ML model is obtained (e.g., after the hyperparameters have been adjusted to achieve a desired level of performance), a third step of collecting the output generated by the trained ML model applied to the third subset (the testing set) may begin. The output generated from the testing set may be compared with the corresponding desired target values to give a final assessment of the trained ML model's accuracy. Other segmentations of the larger data set and/or schemes for using the segments for training one or more ML models are possible.

(74) Backpropagation is an algorithm for training an ML model. Backpropagation is used to adjust (e.g., update) the value of the parameters in the ML model, with the goal of optimizing the objective function. For example, a defined loss function is calculated by forward propagation of an input to obtain an output of the ML model and a comparison of the output value with the target value. Backpropagation calculates a gradient of the loss function with respect to the parameters of the ML model, and a gradient algorithm (e.g., gradient descent) is used to update (e.g., “learn”) the parameters to reduce the loss function. Backpropagation is performed iteratively so that the loss function is converged or minimized. Other techniques for learning the parameters of the ML model can be used. The process of updating (or learning) the parameters over many iterations is referred to as training. Training may be carried out iteratively until a convergence condition is met (e.g., a predefined maximum number of iterations has been performed, or the value outputted by the ML model is sufficiently converged with the desired target value), after which the ML model is considered to be sufficiently trained. The values of the learned parameters can then be fixed and the ML model may be deployed to generate output in real-world applications (also referred to as “inference”).

(75) In some examples, a trained ML model may be fine-tuned, meaning that the values of the learned parameters may be adjusted slightly in order for the ML model to better model a specific task. Fine-tuning of an ML model typically involves further training the ML model on a number of data samples (which may be smaller in number/cardinality than those used to train the model initially) that closely target the specific task. For example, an ML model for generating natural language that has been trained generically on publicly available text corpora may be, e.g., fine-tuned by further training using specific training samples. The specific training samples can be used to generate language in a certain style or in a certain format. For example, the ML model can be trained to generate a blog post having a particular style and structure with a given topic.

(76) Some concepts in ML-based language models are now discussed. It may be noted that, while the term “language model” has been commonly used to refer to an ML-based language model, there could exist non-ML language models. In the present disclosure, the term “language model” can refer to an ML-based language model (e.g., a language model that is implemented using a neural network or other ML architecture), unless stated otherwise. For example, unless stated otherwise, the “language model” encompasses LLMs.

(77) A language model can use a neural network (typically a DNN) to perform natural language processing (NLP) tasks. A language model can be trained to model how words relate to each other in a textual sequence, based on probabilities. A language model may contain hundreds of thousands of learned parameters or, in the case of an LLM, can contain millions or billions of learned parameters or more. As non-limiting examples, a language model can generate text, translate text, summarize text, answer questions, write code (e.g., Python, JavaScript, or other programming languages), classify text (e.g., to identify spam emails), create content for various purposes (e.g., social media content, factual content, or marketing content), or create personalized content for a particular individual or group of individuals. Language models can also be used for chatbots (e.g., virtual assistance).

(78) A type of neural network architecture, referred to as a “transformer,” can be used for language models. For example, the Bidirectional Encoder Representations from Transformers (BERT) model, the Transformer-XL model, and the Generative Pre-trained Transformer (GPT) models are types of transformers. A transformer is a type of neural network architecture that uses self-attention mechanisms in order to generate predicted output based on input data that has some sequential meaning (i.e., the order of the input data is meaningful, which is the case for most text input). Although transformer-based language models are described herein, it should be understood that the present disclosure may be applicable to any ML-based language model, including language models based on other neural network architectures such as recurrent neural network (RNN)-based language models.

(79) FIG. 5 is a block diagram of an example transformer **512**. A transformer is a type of neural network architecture that uses self-attention mechanisms to generate predicted output based on input data that has some sequential meaning (e.g., the order of the input data is meaningful, which is the case for most text input). Self-attention is a mechanism that relates different positions of a single sequence to compute a representation of the same sequence. Although transformer-based language models are described herein, the present disclosure may be applicable to any ML-based language model, including language models based on other neural network architectures such as recurrent neural network (RNN)-based language models.

(80) The transformer **512** includes an encoder **508** (which can include one or more encoder layers/blocks connected in series) and a decoder **510** (which can include one or more decoder layers/blocks connected in series). Generally, the encoder **508** and the decoder **510** each include multiple neural network layers, at least one of which can be a self-attention layer. The parameters of the neural network layers can be referred to as the parameters of the language model.

(81) The transformer **512** can be trained to perform certain functions on a natural language input. Examples of the functions include summarizing existing content, brainstorming ideas, writing a

rough draft, fixing spelling and grammar, and translating content. Summarizing can include extracting key points or themes from an existing content in a high-level summary. Brainstorming ideas can include generating a list of ideas based on provided input. For example, the ML model can generate a list of names for a startup or costumes for an upcoming party. Writing a rough draft can include generating writing in a particular style that could be useful as a starting point for the user's writing. The style can be identified as, e.g., an email, a blog post, a social media post, or a poem. Fixing spelling and grammar can include correcting errors in an existing input text. Translating can include converting an existing input text into a variety of different languages. In some implementations, the transformer **512** is trained to perform certain functions on other input formats than natural language input. For example, the input can include objects, images, audio content, or video content, or a combination thereof.

(82) The transformer **512** can be trained on a text corpus that is labeled (e.g., annotated to indicate verbs, nouns) or unlabeled. LLMs can be trained on a large unlabeled corpus. The term “language model,” as used herein, can include an ML-based language model (e.g., a language model that is implemented using a neural network or other ML architecture), unless stated otherwise. Some LLMs can be trained on a large multi-language, multi-domain corpus to enable the model to be versatile at a variety of language-based tasks such as generative tasks (e.g., generating human-like natural language responses to natural language input).

(83) FIG. 5 illustrates an example of how the transformer **512** can process textual input data. Input to a language model (whether transformer-based or otherwise) typically is in the form of natural language that can be parsed into tokens. The term “token” in the context of language models and NLP has a different meaning from the use of the same term in other contexts such as data security. Tokenization, in the context of language models and NLP, refers to the process of parsing textual input (e.g., a character, a word, a phrase, a sentence, a paragraph) into a sequence of shorter segments that are converted to numerical representations referred to as tokens (or “compute tokens”). Typically, a token can be an integer that corresponds to the index of a text segment (e.g., a word) in a vocabulary dataset. Often, the vocabulary dataset is arranged by frequency of use. Commonly occurring text, such as punctuation, can have a lower vocabulary index in the dataset and thus be represented by a token having a smaller integer value than less commonly occurring text. Tokens frequently correspond to words, with or without white space appended. In some implementations, a token can correspond to a portion of a word.

(84) For example, the word “greater” can be represented by a token for [great] and a second token for [er]. In another example, the text sequence “write a summary” can be parsed into the segments [write], [a], and [summary], each of which can be represented by a respective numerical token. In addition to tokens that are parsed from the textual sequence (e.g., tokens that correspond to words and punctuation), there can also be special tokens to encode non-textual information. For example, a [CLASS] token can be a special token that corresponds to a classification of the textual sequence (e.g., can classify the textual sequence as a list, a paragraph), an [EOT] token can be another special token that indicates the end of the textual sequence, other tokens can provide formatting information, etc.

(85) In FIG. 5, a short sequence of tokens **502** corresponding to the input text is illustrated as input to the transformer **512**. Tokenization of the text sequence into the tokens **502** can be performed by some pre-processing tokenization module such as, for example, a byte-pair encoding tokenizer (the “pre” referring to the tokenization occurring prior to the processing of the tokenized input by the LLM), which is not shown in FIG. 5 for brevity. In general, the token sequence that is inputted to the transformer **512** can be of any length up to a maximum length defined based on the dimensions of the transformer **512**. Each token **502** in the token sequence is converted into an embedding vector **506** (also referred to as “embedding **506**”).

(86) An embedding **506** is a learned numerical representation (such as, for example, a vector) of a token that captures some semantic meaning of the text segment represented by the token **502**. The

embedding **506** represents the text segment corresponding to the token **502** in a way such that embeddings corresponding to semantically related text are closer to each other in a vector space than embeddings corresponding to semantically unrelated text. For example, assuming that the words “write,” “a,” and “summary” each correspond to, respectively, a “write” token, an “a” token, and a “summary” token when tokenized, the embedding **506** corresponding to the “write” token will be closer to another embedding corresponding to the “jot down” token in the vector space as compared to the distance between the embedding **506** corresponding to the “write” token and another embedding corresponding to the “summary” token.

(87) The vector space can be defined by the dimensions and values of the embedding vectors. Various techniques can be used to convert a token **502** to an embedding **506**. For example, another trained ML model can be used to convert the token **502** into an embedding **506**. In particular, another trained ML model can be used to convert the token **502** into an embedding **506** in a way that encodes additional information into the embedding **506** (e.g., a trained ML model can encode positional information about the position of the token **502** in the text sequence into the embedding **506**). In some implementations, the numerical value of the token **502** can be used to look up the corresponding embedding in an embedding matrix **504**, which can be learned during training of the transformer **512**.

(88) The generated embeddings **506** are input into the encoder **508**. The encoder **508** serves to encode the embeddings **506** into feature vectors **514** that represent the latent features of the embeddings **506**. The encoder **508** can encode positional information (i.e., information about the sequence of the input) in the feature vectors **514**. The feature vectors **514** can have very high dimensionality (e.g., on the order of thousands or tens of thousands), with each element in a feature vector **514** corresponding to a respective feature. The numerical weight of each element in a feature vector **514** represents the importance of the corresponding feature. The space of all possible feature vectors **514** that can be generated by the encoder **508** can be referred to as a latent space or feature space.

(89) Conceptually, the decoder **510** is designed to map the features represented by the feature vectors **514** into meaningful output, which can depend on the task that was assigned to the transformer **512**. For example, if the transformer **512** is used for a translation task, the decoder **510** can map the feature vectors **514** into text output in a target language different from the language of the original tokens **502**. Generally, in a generative language model, the decoder **510** serves to decode the feature vectors **514** into a sequence of tokens. The decoder **510** can generate output tokens **516** one by one. Each output token **516** can be fed back as input to the decoder **510** in order to generate the next output token **516**. By feeding back the generated output and applying self-attention, the decoder **510** can generate a sequence of output tokens **516** that has sequential meaning (e.g., the resulting output text sequence is understandable as a sentence and obeys grammatical rules). The decoder **510** can generate output tokens **516** until a special [EOT] token (indicating the end of the text) is generated. The resulting sequence of output tokens **516** can then be converted to a text sequence in post-processing. For example, each output token **516** can be an integer number that corresponds to a vocabulary index. By looking up the text segment using the vocabulary index, the text segment corresponding to each output token **516** can be retrieved, the text segments can be concatenated together, and the final output text sequence can be obtained.

(90) In some implementations, the input provided to the transformer **512** includes instructions to perform a function on an existing text. The output can include, for example, a modified version of the input text and instructions to modify the text. The modification can include summarizing, translating, correcting grammar or spelling, changing the style of the input text, lengthening or shortening the text, or changing the format of the text (e.g., adding bullet points or checkboxes). As an example, the input text can include meeting notes prepared by a user and the output can include a high-level summary of the meeting notes. In other examples, the input provided to the transformer includes a question or a request to generate text. The output can include a response to

the question, text associated with the request, or a list of ideas associated with the request. For example, the input can include the question “What is the weather like in San Francisco?” and the output can include a description of the weather in San Francisco. As another example, the input can include a request to brainstorm names for a flower shop and the output can include a list of relevant names.

(91) Although a general transformer architecture for a language model and its theory of operation have been described above, this is not intended to be limiting. Existing language models include language models that are based only on the encoder of the transformer or only on the decoder of the transformer. An encoder-only language model encodes the input text sequence into feature vectors that can then be further processed by a task-specific layer (e.g., a classification layer). BERT is an example of a language model that can be considered to be an encoder-only language model. A decoder-only language model accepts embeddings as input and can use auto-regression to generate an output text sequence. Transformer-XL and GPT-type models can be language models that are considered to be decoder-only language models.

(92) Because GPT-type language models tend to have a large number of parameters, these language models can be considered LLMs. An example of a GPT-type LLM is GPT-3. GPT-3 is a type of GPT language model that has been trained (in an unsupervised manner) on a large corpus derived from documents available online to the public. GPT-3 has a very large number of learned parameters (on the order of hundreds of billions), can accept a large number of tokens as input (e.g., up to 2,048 input tokens), and is able to generate a large number of tokens as output (e.g., up to 2,048 tokens). GPT-3 has been trained as a generative model, meaning that it can process input text sequences to predictively generate a meaningful output text sequence. ChatGPT is built on top of a GPT-type LLM and has been fine-tuned with training datasets based on text-based chats (e.g., chatbot conversations). ChatGPT is designed for processing natural language, receiving chat-like inputs, and generating chat-like outputs.

(93) A computer system can access a remote language model (e.g., a cloud-based language model), such as ChatGPT or GPT-3, via a software interface (e.g., an API). Additionally or alternatively, such a remote language model can be accessed via a network such as the Internet. In some implementations, such as, for example, potentially in the case of a cloud-based language model, a remote language model can be hosted by a computer system that can include a plurality of cooperating (e.g., cooperating via a network) computer systems that can be in, for example, a distributed arrangement. Notably, a remote language model can employ multiple processors (e.g., hardware processors such as, for example, processors of cooperating computer systems). Indeed, processing of inputs by an LLM can be computationally expensive/can involve a large number of operations (e.g., many instructions can be executed/large data structures can be accessed from memory), and providing output in a required timeframe (e.g., real time or near real time) can require the use of a plurality of processors/cooperating computing devices as discussed above.

(94) Inputs to an LLM can be referred to as a prompt, which is a natural language input that includes instructions to the LLM to generate a desired output. A computer system can generate a prompt that is provided as input to the LLM via an API (e.g., the API 328 in FIG. 3). As described above, the prompt can optionally be processed or pre-processed into a token sequence prior to being provided as input to the LLM via its API. A prompt can include one or more examples of the desired output, which provides the LLM with additional information to enable the LLM to generate output according to the desired output. Additionally or alternatively, the examples included in a prompt can provide inputs (e.g., example inputs) corresponding to/as can be expected to result in the desired outputs provided. A one-shot prompt refers to a prompt that includes one example, and a few-shot prompt refers to a prompt that includes multiple examples. A prompt that includes no examples can be referred to as a zero-shot prompt.

(95) AI Assistant Workflow

(96) FIG. 6 is a flowchart illustrating a process 600 for controlling an environment using an

artificial intelligence assistant, according to some implementations. The process **600** can be performed by a computing system, such as the assistant **120** described above or a computing system that implements the assistant **120**. Other implementations of the process **600** include additional, fewer, or different steps, or perform the steps in different orders.

(97) As shown in FIG. **6**, the computer system receives, at **602**, an instruction to perform a task within an environment that is communicatively coupled to the computer system, such as the controlled environment **110**. The instruction can be a natural language instruction that is input by a user of the computer system or the environment. In an example, a user interacts with the computer system to perform a task within the platform **300** described with respect to FIG. **3**. A user instruction can be input via any of a variety of sources, such as a chat-like interface shown in FIG. **2B** or a text entry box displayed within a page of the environment as shown in FIG. **2D**.

(98) At **604**, the computer system generates a computer-readable input based on the received instructions. The computer-readable input can include a context of the environment, as well as a computer-readable form of the user's natural language instructions.

(99) At **606**, the computer system sends the computer-readable input to a large language model (LLM) to cause the LLM to generate a set of computer program code, such as XML or JavaScript, to perform the task requested by the user. The LLM can be trained using pairs of user instructions and code to perform the instructions, such that the LLM is configured to process the computer-readable input and to generate computer program code in response. Depending on the task that is to be performed or the structure of the environment, the computer system can cause the LLM to generate different types of computer program code. The computer system can validate the output of the LLM to ensure that the code is correct, such as verifying that it employs real APIs, functions, and syntax within the appropriate coding language.

(100) At **608**, after the LLM has returned the generated computer program code, the computer system executes the code to perform the task in the environment.

(101) Using the process illustrated in FIG. **6**, the computer system can iteratively add computer-readable inputs and LLM-produced code to a transcript that enables the computer system to perform tasks in the environment. After performing a first task at **608**, the computer system can receive subsequent natural language inputs from users. Based on these subsequent inputs, the computer system can correct previous task results (e.g., by instructing the LLM to correct previously generated code or to output new code based on updated contexts or task parameters) or perform additional tasks in the environment.

(102) XML Interpreter

(103) As discussed with respect to FIG. **6**, some implementations of a workflow of interactions between an AI assistant (such as the assistant **120**) and an LLM include causing the LLM to generate XML to perform various tasks. However, the XML produced by LLM may sometimes fail to perform the desired task, whether due to hallucinations by the LLM, errors in the prompts generated by the assistant, modifications to or ambiguities in the environment, or other inherent or exogenous factors. To ensure that tasks can be performed correctly, some implementations of the assistant **120** therefore employ a process to validate the XML that is produced by the LLM **130**.

(104) FIG. **7** is a flowchart illustrating a process **700** for implementing an XML interpreter to validate instructions that are written by or with the assistance of a large language model, according to some implementations. The process **700** can be performed by a computer system such as the assistant **120** described above. Other implementations of the process **700** include additional, fewer, or different steps, or perform the steps in different orders.

(105) At **702**, the computer system uses an LLM to generate a first set of XML instructions associated with performing a first task in an environment coupled to the computing system. Like implementations described above, the environment in which the task is performed can be the data management platform, a collaboration platform, or any other physical or virtual environment whose state or content can be remotely read and/or modified by a computing system. The first task can

include any action to retrieve a current state of the environment or to modify a current state of the environment, such as observing a value stored in the data management platform or writing content to a chat thread within the collaboration platform. In some cases, the computer system causes the LLM to generate the first set of XML instructions in response to a natural language input from a user to perform the first task or to perform a series of tasks that includes the first task. Alternatively, the computer system can trigger generation of the first set of XML in response to another task performed by the computer system.

(106) After receiving a first set of XML instructions from the LLM, the computer system executes the first set of instructions (at **704**) and observes a result of the execution (at **706**). The computer system can evaluate, for example, whether the instructions are executable or whether the instructions cannot be executed due to syntax errors, hallucinated API calls, or the like. If the instructions cannot be executed, the computer system can further process the first set of instructions to identify a cause of the error. Some implementations of the system can perform some automated evaluation of the first set of instructions, such as verifying that each tag is closed. Alternatively, the system can send a non-executable line of XML instructions to an LLM, optionally a different LLM than the LLM used to initially produce the XML, to ask the LLM to identify any errors in the non-executable line. Identifications of lines with errors can also be output to a user of the computer system for review by the user.

(107) If the first set of XML instructions can be executed, the computer system can evaluate whether execution of the instructions yields an expected result at **708**. For example, if the instruction is to write content to a specified location within a structured digital environment, the computer system evaluates whether the specified location can be found in the environment, whether the content can be identified or generated, and—after performing the write operation—observing whether the correct content was written to the correct location. For example, the computer system may be unable to find a location within the environment if there is no location in the environment that has the name given in the first set of XML instructions. Content may be unidentifiable if the content that is to be written depends on another content item that cannot be located or if the XML instructions fail to properly handle a prompt back into the LLM to generate the content. Similarly, if the instruction is to read a value from the environment, the computer system can evaluate whether execution of the instructions returns a read value or whether the read operation returns a null value.

(108) If execution of the first set of XML instructions does not return an expected result, the computer system can use the LLM to modify the first set of XML instructions, at **710**. Depending on the error detected, the system may provide the original set of instructions back to the LLM, with a request to change an aspect of the XML that is produced. For other types of errors, the system can the LLM to generate new instructions or can ask a different LLM to generate the instructions.

(109) The computer system can repeat the operations at **704-710** until a set of instructions that correctly perform the first task has been generated.

(110) Once the first set of XML instructions have been determined to be executable to produce an expected result, the computer system executes the first set of XML, at **712**, to perform the first task in the environment. For example, the computer system can write content to the environment or delete content from the environment according to the instructions for the first task.

(111) Finally, at **714**, the computer system uses the LLM to generate a second set of XML instructions to perform a second task in the environment. The second set of XML instructions can be generated concurrently with the first set of XML instructions or after the first set of instructions have been determined to produce an expected result, for example. For the second set of instructions, the computer system can use a similar process as that described above to validate the second set of instructions.

(112) Computer System

(113) FIG. **8** is a block diagram that illustrates an example of a computer system **800** in which at

least some operations described herein can be implemented. As shown, the computer system **800** can include: one or more processors **802**, main memory **806**, non-volatile memory **810**, a network interface device **812**, a display device **818**, an input/output device **820**, a control device **822** (e.g., keyboard and pointing device), a drive unit **824** that includes a machine readable (storage) medium **826**, and a signal generation device **830** that are communicatively connected to a bus **816**. The bus **816** represents one or more physical buses and/or point-to-point connections that are connected by appropriate bridges, adapters, or controllers. Various common components (e.g., cache memory) are omitted from FIG. **8** for brevity. Instead, the computer system **800** is intended to illustrate a hardware device on which components illustrated or described relative to the examples of the figures and any other components described in this specification can be implemented.

(114) The computer system **800** can take any suitable physical form. For example, the computer system **800** can share a similar architecture as that of a server computer, personal computer (PC), tablet computer, mobile telephone, wearable electronic device, network-connected (“smart”) device (e.g., a television or home assistant device), AR/VR system (e.g., head-mounted display), or any electronic device capable of executing a set of instructions that specify action(s) to be taken by the computer system **800**. In some implementations, the computer system **800** can be an embedded computer system, a system-on-chip (SOC), a single-board computer (SBC) system, or a distributed system such as a mesh of computer systems or include one or more cloud components in one or more networks. Where appropriate, one or more computer systems **800** can perform operations in real time, near real time, or in batch mode.

(115) The network interface device **812** enables the computer system **800** to mediate data in a network **814** with an entity that is external to the computer system **800** through any communication protocol supported by the computer system **800** and the external entity. Examples of the network interface device **812** include a network adapter card, a wireless network interface card, a router, an access point, a wireless router, a switch, a multilayer switch, a protocol converter, a gateway, a bridge, bridge router, a hub, a digital media receiver, and/or a repeater, as well as all wireless elements noted herein.

(116) The memory (e.g., main memory **806**, non-volatile memory **810**, machine-readable medium **826**) can be local, remote, or distributed. Although shown as a single medium, the machine-readable medium **826** can include multiple media (e.g., a centralized/distributed database and/or associated caches and servers) that store one or more sets of instructions **828**. The machine-readable medium **826** can include any medium that is capable of storing, encoding, or carrying a set of instructions for execution by the computer system **800**. The machine-readable medium **826** can be non-transitory or comprise a non-transitory device. In this context, a non-transitory storage medium can include a device that is tangible, meaning that the device has a concrete physical form, although the device can change its physical state. Thus, for example, non-transitory refers to a device remaining tangible despite this change in state.

(117) Although implementations have been described in the context of fully functioning computing devices, the various examples are capable of being distributed as a program product in a variety of forms. Examples of machine-readable storage media, machine-readable media, or computer-readable media include recordable-type media such as volatile and non-volatile memory devices **810**, removable flash memory, hard disk drives, optical disks, and transmission-type media such as digital and analog communication links.

(118) In general, the routines executed to implement examples herein can be implemented as part of an operating system or a specific application, component, program, object, module, or sequence of instructions (collectively referred to as “computer programs”). The computer programs typically comprise one or more instructions (e.g., instructions **804**, **808**, **828**) set at various times in various memory and storage devices in computing device(s). When read and executed by the processor **802**, the instruction(s) cause the computer system **800** to perform operations to execute elements involving the various aspects of the disclosure.

(119) Remarks

(120) The terms “example,” “embodiment,” and “implementation” are used interchangeably. For example, references to “one example” or “an example” in the disclosure can be, but not necessarily are, references to the same implementation; and such references mean at least one of the implementations. The appearances of the phrase “in one example” are not necessarily all referring to the same example, nor are separate or alternative examples mutually exclusive of other examples. A feature, structure, or characteristic described in connection with an example can be included in another example of the disclosure. Moreover, various features are described that can be exhibited by some examples and not by others. Similarly, various requirements are described that can be requirements for some examples but not other examples.

(121) The terminology used herein should be interpreted in its broadest reasonable manner, even though it is being used in conjunction with certain specific examples of the invention. The terms used in the disclosure generally have their ordinary meanings in the relevant technical art, within the context of the disclosure, and in the specific context where each term is used. A recital of alternative language or synonyms does not exclude the use of other synonyms. Special significance should not be placed upon whether or not a term is elaborated or discussed herein. The use of highlighting has no influence on the scope and meaning of a term. Further, it will be appreciated that the same thing can be said in more than one way.

(122) Unless the context clearly requires otherwise, throughout the description and the claims, the words “comprise,” “comprising,” and the like are to be construed in an inclusive sense, as opposed to an exclusive or exhaustive sense; that is to say, in the sense of “including, but not limited to.” As used herein, the terms “connected,” “coupled,” or any variant thereof means any connection or coupling, either direct or indirect, between two or more elements; the coupling or connection between the elements can be physical, logical, or a combination thereof. Additionally, the words “herein,” “above,” “below,” and words of similar import can refer to this application as a whole and not to any particular portions of this application. Where context permits, words in the Detailed Description above using the singular or plural number may also include the plural or singular number respectively. The word “or” in reference to a list of two or more items covers all of the following interpretations of the word: any of the items in the list, all of the items in the list, and any combination of the items in the list. The term “module” refers broadly to software components, firmware components, and/or hardware components.

(123) While specific examples of technology are described above for illustrative purposes, various equivalent modifications are possible within the scope of the invention, as those skilled in the relevant art will recognize. For example, while processes or blocks are presented in a given order, alternative implementations can perform routines having steps, or employ systems having blocks, in a different order, and some processes or blocks may be deleted, moved, added, subdivided, combined, and/or modified to provide alternative or sub-combinations. Each of these processes or blocks can be implemented in a variety of different ways. Also, while processes or blocks are at times shown as being performed in series, these processes or blocks can instead be performed or implemented in parallel, or can be performed at different times. Further, any specific numbers noted herein are only examples such that alternative implementations can employ differing values or ranges.

(124) Details of the disclosed implementations can vary considerably in specific implementations while still being encompassed by the disclosed teachings. As noted above, particular terminology used when describing features or aspects of the invention should not be taken to imply that the terminology is being redefined herein to be restricted to any specific characteristics, features, or aspects of the invention with which that terminology is associated. In general, the terms used in the following claims should not be construed to limit the invention to the specific examples disclosed herein, unless the Detailed Description above explicitly defines such terms. Accordingly, the actual scope of the invention encompasses not only the disclosed examples but also all equivalent ways of

practicing or implementing the invention under the claims. Some alternative implementations can include additional elements to those implementations described above or include fewer elements. (125) Any patents and applications and other references noted above, and any that may be listed in accompanying filing papers, are incorporated herein by reference in their entireties, except for any subject matter disclaimers or disavowals, and except to the extent that the incorporated material is inconsistent with the express disclosure herein, in which case the language in this disclosure controls. Aspects of the invention can be modified to employ the systems, functions, and concepts of the various references described above to provide yet further implementations of the invention. (126) To reduce the number of claims, certain implementations are presented below in certain claim forms, but the applicant contemplates various aspects of an invention in other forms. For example, aspects of a claim can be recited in a means-plus-function form or in other forms, such as being embodied in a computer-readable medium. A claim intended to be interpreted as a mean-plus-function claim will use the words “means for.” However, the use of the term “for” in any other context is not intended to invoke a similar interpretation. The applicant reserves the right to pursue such additional claim forms in either this application or in a continuing application.

Claims

1. A non-transitory, computer-readable storage medium comprising instructions recorded thereon, wherein the instructions, when executed by at least one data processor of a system, cause the system to: use a large language model (LLM) to generate a first set of extensible markup language (XML) instructions to perform a first task in an environment communicatively coupled to the system; execute the first set of XML instructions; observe a result of the execution; and based on the observed result of the execution: use the LLM to generate a second set of XML instructions to perform a second task in the environment; or detect an error in the first set of XML instructions and, in response to detecting the error, use the LLM to modify the first set of XML instructions to correct the error.
2. The non-transitory computer-readable storage medium of claim 1, wherein the instructions when executed further cause the system to: after observing the result of the execution and using the LLM to modify the first set of XML instructions to correct the error, executing the modified first set of XML instructions to output a task result for the first task to the environment.
3. The non-transitory computer-readable storage medium of claim 1, wherein the instructions when executed cause the system to use the LLM to generate the second set of XML instructions in response to detecting no error in the first set of XML instructions.
4. The non-transitory computer-readable storage medium of claim 1, wherein the first task includes writing content to a specified location within the environment, and wherein observing the result of the execution comprises: parsing the first set of XML instructions to determine an identification of the specified location; and observing the environment to locate the specified location.
5. The non-transitory computer-readable storage medium of claim 4, wherein the error is detected when the specified location is not located in the environment.
6. The non-transitory computer-readable storage medium of claim 1, wherein the first task includes writing content to the environment, and wherein observing the result of the execution comprises: parsing the first set of XML instructions to determine whether the content to be written can be identified.
7. The non-transitory computer-readable storage medium of claim 6, wherein the error is detected when the first set of XML instructions cannot be executed to identify the content to be written.
8. The non-transitory computer-readable storage medium of claim 1, wherein the first task includes reading content from the environment, and wherein executing the first set of XML instructions comprises performing a read operation to read the content from the environment.
9. The non-transitory computer-readable storage medium of claim 8, wherein detecting the error

comprises detecting the read operation returns a null value.

10. The non-transitory computer-readable storage medium of claim 1, wherein observing the result of the execution comprises detecting a syntax error in the first set of XML instructions.

11. The non-transitory computer-readable storage medium of claim 1, wherein the instructions when executed cause the system to use the LLM to generate the first set of XML instructions in response to a natural language input received from a user in association with the environment.

12. The non-transitory computer-readable storage medium of claim 1, wherein the instructions when executed cause the system to use the LLM to generate the first set of XML instructions in response to execution of a third set of XML instructions to perform a third task in the environment.

13. A system comprising: at least one hardware processor; and at least one non-transitory memory storing instructions, which, when executed by the at least one hardware processor, cause the system to: use a large language model (LLM) to generate a first set of extensible markup language (XML) instructions to perform a first task in an environment communicatively coupled to the system; execute the first set of XML instructions; observe a result of the execution; and based on the observed result of the execution: use the LLM to generate a second set of XML instructions to perform a second task in the environment; or detect an error in the first set of XML instructions and, in response to detecting the error, use the LLM to modify the first set of XML instructions to correct the error.

14. The system of claim 13, wherein the instructions when executed further cause the system to: after observing the result of the execution and using the LLM to modify the first set of XML instructions to correct the error, executing the modified first set of XML instructions to output a task result for the first task to the environment.

15. The system of claim 13, wherein the first task includes writing content to a specified location within the environment, and wherein observing the result of the execution comprises: parsing the first set of XML instructions to determine an identification of the specified location; and observing the environment to locate the specified location.

16. The system of claim 13, wherein observing the result of the execution comprises detecting a syntax error in the first set of XML instructions.

17. A method comprising: causing a computer system to use a large language model (LLM) to generate a first set of extensible markup language (XML) instructions to perform a first task in an environment communicatively coupled to the computer system; executing, by the computer system, the first set of XML instructions; observing, by the computer system, a result of the execution; and based on the observed result of the execution: causing the computer system to use the LLM to generate a second set of XML instructions to perform a second task in the environment; or detecting, by the computer system, an error in the first set of XML instructions and, in response to detecting the error, use the LLM to modify the first set of XML instructions to correct the error.

18. The method of claim 17, further comprising: after observing the result of the execution and using the LLM to modify the first set of XML instructions to correct the error, executing, by the computer system, the modified first set of XML instructions to output a task result for the first task to the environment.

19. The method of claim 17, wherein the first task includes writing content to a specified location within the environment, and wherein observing the result of the execution comprises: parsing the first set of XML instructions to determine an identification of the specified location; and observing the environment to locate the specified location.

20. The method of claim 17, wherein observing the result of the execution comprises detecting a syntax error in the first set of XML instructions.
