

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent
Kind Code
Date of Patent
Inventor(s)

12405738
B2
September 02, 2025
Charles; Peniel et al.

Synchronized mapping for cache-aware load balancing across multiple host devices

Abstract

An apparatus in one embodiment comprises at least one processing device configured to store a mapping for each of a plurality of logical storage devices of a storage system, the mapping for a given one of the logical storage devices assigning different ranges of logical block addresses of that logical storage device to respective different cache entities of the storage system. The at least one processing device is further configured to receive at least one mapping-related communication from at least one of first and second host devices that have shared access to the given logical storage device, to modify the stored mapping for the given logical storage device based at least in part on the at least one received mapping-related communication, and to send at least one mapping-related communication to at least one of the first and second host devices based at least in part on the modified stored mapping.

Inventors: Charles; Peniel (Bangalore, IN), Nagaraj; Venkatesh Doddamane (Bangalore, IN), Anchi; Amit Pundalik (Bangalore, IN)

Applicant: Dell Products L.P. (Round Rock, TX)

Family ID: 89324002

Assignee: Dell Products L.P. (Round Rock, TX)

Appl. No.: 17/847300

Filed: June 23, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20230418492 A1	Dec. 28, 2023

Publication Classification

Int. Cl.: G06F3/00 (20060101); **G06F3/06** (20060101); **G06F12/0815** (20160101)

U.S. Cl.:

CPC **G06F3/064** (20130101); **G06F3/0604** (20130101); **G06F3/0635** (20130101);
G06F3/0659 (20130101); **G06F3/0683** (20130101); **G06F12/0815** (20130101);

Field of Classification Search

CPC: G06F (12/0815); G06F (3/0604); G06F (3/064); G06F (3/0635); G06F (3/0659)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6567397	12/2002	Campana, Jr. et al.	N/A	N/A
6687746	12/2003	Shuster et al.	N/A	N/A
6697875	12/2003	Wilson	N/A	N/A
7275103	12/2006	Thrasher et al.	N/A	N/A
7454437	12/2007	Lavallee et al.	N/A	N/A
7617292	12/2008	Moore et al.	N/A	N/A
7668981	12/2009	Nagineni et al.	N/A	N/A
7770053	12/2009	Bappe et al.	N/A	N/A
7809912	12/2009	Raizen et al.	N/A	N/A
7818428	12/2009	Lavallee et al.	N/A	N/A
7890664	12/2010	Tao et al.	N/A	N/A
7904681	12/2010	Bappe et al.	N/A	N/A
7925872	12/2010	Lai et al.	N/A	N/A
8250256	12/2011	Ghosalkar et al.	N/A	N/A
8285825	12/2011	Nagaraj et al.	N/A	N/A
8825919	12/2013	Lim et al.	N/A	N/A
8832334	12/2013	Okita	N/A	N/A
8874746	12/2013	Gonzalez	N/A	N/A
8972657	12/2014	Armangau et al.	N/A	N/A
9026694	12/2014	Davidson et al.	N/A	N/A
9201803	12/2014	Derbeko et al.	N/A	N/A
9400611	12/2015	Raizen	N/A	N/A
9430368	12/2015	Derbeko et al.	N/A	N/A
9594780	12/2016	Esposito et al.	N/A	N/A
9647933	12/2016	Tawri et al.	N/A	N/A
9672160	12/2016	Derbeko et al.	N/A	N/A
9733846	12/2016	Wigmore et al.	N/A	N/A
9778852	12/2016	Marshak et al.	N/A	N/A
10289325	12/2018	Bono	N/A	N/A
10353714	12/2018	Gokam et al.	N/A	N/A
10439878	12/2018	Tah et al.	N/A	N/A
10474367	12/2018	Mallick et al.	N/A	N/A
10476960	12/2018	Rao et al.	N/A	N/A
10521369	12/2018	Mallick et al.	N/A	N/A

10606496	12/2019	Mallick et al.	N/A	N/A
10637917	12/2019	Mallick et al.	N/A	N/A
10652206	12/2019	Pusalkar et al.	N/A	N/A
10754572	12/2019	Kumar et al.	N/A	N/A
10757189	12/2019	Mallick et al.	N/A	N/A
10764371	12/2019	Rao et al.	N/A	N/A
10789006	12/2019	Gokam et al.	N/A	N/A
10817181	12/2019	Mallick et al.	N/A	N/A
10838648	12/2019	Sharma et al.	N/A	N/A
10880217	12/2019	Mallick et al.	N/A	N/A
10884935	12/2020	Doddaiah	N/A	N/A
10911402	12/2020	Pusalkar et al.	N/A	N/A
11086785	12/2020	Anchi et al.	N/A	N/A
11151071	12/2020	Patel	N/A	G06F 13/4022
2002/0023151	12/2001	Iwatani	N/A	N/A
2002/0103923	12/2001	Cherian et al.	N/A	N/A
2004/0010563	12/2003	Forte et al.	N/A	N/A
2006/0026346	12/2005	Kadoiri et al.	N/A	N/A
2006/0277383	12/2005	Hayden et al.	N/A	N/A
2007/0174849	12/2006	Cheung et al.	N/A	N/A
2008/0043973	12/2007	Lai et al.	N/A	N/A
2008/0201458	12/2007	Salli	N/A	N/A
2008/0301332	12/2007	Butler et al.	N/A	N/A
2009/0210634	12/2008	Tsuruta	711/E12.001	G06F 11/1004
2009/0259749	12/2008	Barrett et al.	N/A	N/A
2009/0292834	12/2008	Neemidge et al.	N/A	N/A
2010/0251267	12/2009	Zwisler et al.	N/A	N/A
2010/0313063	12/2009	Venkataraja et al.	N/A	N/A
2011/0029730	12/2010	Durocher et al.	N/A	N/A
2011/0197027	12/2010	Balasubramanian et al.	N/A	N/A
2011/0296230	12/2010	Chen et al.	N/A	N/A
2012/0102369	12/2011	Hiltunen et al.	N/A	N/A
2012/0246345	12/2011	Contreras et al.	N/A	N/A
2013/0117766	12/2012	Bax et al.	N/A	N/A
2013/0290571	12/2012	Rizzo et al.	N/A	N/A
2013/0339551	12/2012	Flanagan et al.	N/A	N/A
2014/0105068	12/2013	Xu	N/A	N/A
2015/0169233	12/2014	Anderson et al.	N/A	N/A
2015/0222705	12/2014	Stephens	N/A	N/A
2015/0242134	12/2014	Takada et al.	N/A	N/A
2016/0055101	12/2015	Gerhart	713/193	H04L 9/0894
2016/0092136	12/2015	Balakrishnan et al.	N/A	N/A
2016/0092364	12/2015	Feng et al.	N/A	N/A
2016/0117113	12/2015	Li et al.	N/A	N/A
2016/0335003	12/2015	Ahmed et al.	N/A	N/A
2017/0235507	12/2016	Sinha et al.	N/A	N/A

2018/0189635	12/2017	Olarig et al.	N/A	N/A
2018/0253256	12/2017	Bharadwaj	N/A	N/A
2018/0317101	12/2017	Koue	N/A	N/A
2019/0095299	12/2018	Liu et al.	N/A	N/A
2019/0108888	12/2018	Sarkar et al.	N/A	N/A
2019/0334987	12/2018	Mallick et al.	N/A	N/A
2020/0021653	12/2019	Rao et al.	N/A	N/A
2020/0097203	12/2019	Mallick et al.	N/A	N/A
2020/0106698	12/2019	Rao et al.	N/A	N/A
2020/0110552	12/2019	Kumar et al.	N/A	N/A
2020/0112608	12/2019	Patel et al.	N/A	N/A
2020/0160918	12/2019	Byun	N/A	G06F 3/0659
2020/0192588	12/2019	Kumar et al.	N/A	N/A
2020/0204475	12/2019	Mallick et al.	N/A	N/A
2020/0204495	12/2019	Mallick et al.	N/A	N/A
2020/0213274	12/2019	Pusalkar et al.	N/A	N/A
2020/0241890	12/2019	Mallick et al.	N/A	N/A
2020/0314218	12/2019	Kumar et al.	N/A	N/A
2020/0348860	12/2019	Mallick et al.	N/A	N/A
2020/0348861	12/2019	Marappan et al.	N/A	N/A
2020/0348869	12/2019	Gokam	N/A	N/A
2020/0349094	12/2019	Smith et al.	N/A	N/A
2020/0363985	12/2019	Gokam et al.	N/A	N/A
2020/0372401	12/2019	Mallick et al.	N/A	N/A
2021/0019054	12/2020	Anchi et al.	N/A	N/A
2021/0026551	12/2020	Tidke et al.	N/A	N/A
2021/0026650	12/2020	Rao et al.	N/A	N/A
2021/0103506	12/2020	Marappan	N/A	H04L 67/55
2021/0157502	12/2020	Rao et al.	N/A	N/A
2021/0181965	12/2020	Anchi et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
103677927	12/2016	CN	N/A
1117028	12/2000	EP	N/A
2667569	12/2012	EP	N/A

OTHER PUBLICATIONS

International Search Report and Written Opinion of PCT/US2019/052549 dated Dec. 4, 2019, 13 pages. cited by applicant

International Search Report and Written Opinion of PCT/US2019/053204 dated Dec. 16, 2019, 40 pages. cited by applicant

International Search Report and Written Opinion of PCT/US2019/053473 dated Dec. 19, 2019, 16 pages. cited by applicant

International Search Report and Written Opinion of PCT/US2019/067144 dated May 4, 2020, 26 pages. cited by applicant

Kris Piepho, "Dell EMC SC Series Storage: Microsoft Multipath I/O," Dell EMC Best Practices, Jan. 2017, 57 pages. cited by applicant

NVM Express, “NVM Express, Revision 1.3,” NVM Express, May 1, 2017, 282 pages. cited by applicant

VMware, “Multipathing Configuration for Software iSCSI Using Port Binding,” Technical White Paper, Apr. 25, 2012, 15 pages. cited by applicant

Dell EMC, “Dell EMC SC Series Storage: Microsoft Multipath I/O,” Dell EMC Engineering, Jun. 2017, 56 pages. cited by applicant

EMC, “EMC PowerPath and PowerPath/VE Family for Windows,” Installation and Administration Guide, Oct. 2018, 102 pages. cited by applicant

EMC, “EMC Powerpath Load Balancing and Failover”, Comparison with native MPIO operating system solutions, Feb. 2011, 28 pages. cited by applicant

Dell EMC, “PowerMax OS,” Dell EMC PowerMax Family Product Guide, May 2019, 192 pages. cited by applicant

Dell EMC, “Dell EMC SC Series Storage and Microsoft Multipath I/O,” CML 1004, Jul. 2018, 36 pages. cited by applicant

Dell EMC, “Dell EMC Unity: Virtualization Integration,” Technical White Paper, Oct. 2019, 39 pages. cited by applicant

Dell EMC, “Dell EMC PowerMax: iSCSI Implementation for Dell EMC Storage Arrays Running PowerMaxOS,” Technical White Paper, Sep. 2019, 35 pages. cited by applicant

VMware Knowledge Base, “8GB Emulex HBAs May Not React to RSCN Events Properly on Older Firmware Versions,” <https://kb.vmware.com/s/article/2005235>, Sep. 20, 2013, 1 page. cited by applicant

NetApp Knowledge Base, “SAN Host Does Not Failover I/O to Alternate Paths When One or More Paths Are Removed and an RSCN Is Not Received,” https://kb.netapp.com/app/answers/answer_view/a_id/1074454/~san-host-does-not-failover-i%2Fo-to-alternate-paths-when-one-or-more-paths-are, Nov. 11, 2018, 3 pages. cited by applicant

RedHat, “Device-Mapper-Multipath on RHEL5 Experiences Excessive Delay in Detecting a Lost Path from a Storage Failure that Produces No RSCN or Loop/Link Error,” <https://access.redhat.com/solutions/23576>, Mar. 21, 2014, 2 pages. cited by applicant

E. Goggin et al., “Linux Multipathing,” Proceedings of the Linux Symposium, vol. 1, Jul. 2005, pp. 147-167. cited by applicant

K. Ueda et al., “Request-based Device-mapper Multipath and Dynamic Load Balancing,” Proceedings of the Linux Symposium, vol. 2, Jun. 2007, pp. 235-243. cited by applicant

B. Tulman, “In-Band and out-of-Band Network Management,” <http://www.learncomputer.com/in-band-out-of-band-network-management/>, May 18, 2010, 2 pages. cited by applicant

NVM Express, “NVM Express Base Specification, Revision 2.0b,” NVM Express, Jan. 6, 2022, 455 pages. cited by applicant

U.S. Appl. No. 17/383,672 filed in the name of Amit Pundalik Anchi et al. filed Jul. 23, 2021, and entitled “Multi-Path Layer Configured to Access Storage-Side Performance Metrics for Load Balancing Policy Control.” cited by applicant

U.S. Appl. No. 17/503,562 filed in the name of Amit Pundalik Anchi et al. filed Oct. 18, 2021, and entitled “Dynamic Chunk Size Adjustment for Cache-Aware Load Balancing.” cited by applicant

Primary Examiner: Bragdon; Reginald G

Assistant Examiner: Westbrook; Michael L

Attorney, Agent or Firm: Ryan, Mason & Lewis, LLP

Background/Summary

FIELD

(1) The field relates generally to information processing systems, and more particularly to storage in information processing systems.

BACKGROUND

(2) Storage arrays and other types of storage systems are often shared by multiple host devices over a network. Applications running on the host devices each include one or more processes that perform the application functionality. The processes issue input-output (IO) operations directed to particular logical storage volumes or other logical storage devices, for delivery by the host devices over selected paths to storage ports of the storage system. Different ones of the host devices can run different applications with varying workloads and associated IO patterns. Such host devices also generate additional IO operations in performing various data services such as migration and replication. In these and other situations involving processing of IO requests received from a host device, a storage system may utilize one or more local caches in processing the received IO requests. For example, write requests received from a given host device may be initially stored in one or more write caches of the storage system, and then acknowledged back to the host device, before being destaged to a destination storage device of the storage system. Performance of the storage system in these and other situations that involve processing IO requests received from one or more host devices can be adversely impacted by cache-related issues of the storage system.

SUMMARY

(3) Illustrative embodiments provide synchronized mapping of logical storage volume address ranges to cache entities of a storage system across multiple host devices for storage cache aware load balancing. The host devices illustratively communicate with a storage array or other type of storage system via a storage area network (SAN) or other type of network. The host devices illustratively comprise a multi-path layer that includes multi-path input-output (MPIO) drivers configured to process IO operations of respective ones of the host devices.

(4) These and other illustrative embodiments implement a storage cache aware load balancing policy with synchronized mapping across multiple host devices that can avoid the above-noted adverse impacts of cache-related issues of the storage system.

(5) In one embodiment, an apparatus comprises at least one processing device that includes a processor and a memory, with the processor being coupled to the memory. The at least one processing device is configured to store a mapping for each of a plurality of logical storage devices of a storage system, with the mapping for a given one of the logical storage devices assigning different ranges of logical block addresses (LBAs) of the given logical storage device to respective different cache entities of the storage system.

(6) The at least one processing device is further configured to receive at least one mapping-related communication from at least one of first and second host devices that have shared access to the given logical storage device, to modify the stored mapping for the given logical storage device based at least in part on the at least one received mapping-related communication, and to send at least one mapping-related communication to at least one of the first and second host devices based at least in part on the modified stored mapping.

(7) The at least one processing device illustratively comprises an orchestrator device that is separate from the first and second host devices. In some embodiments, the orchestrator device more particularly comprises an MPIO management station coupled to each of the first and second host devices and configured to interact with first and second MPIO drivers of the respective first and second host devices.

(8) The first and second host devices are illustratively part of a host device cluster that in some

embodiments can include a potentially very large number of host devices, each in communication with the orchestrator device. The above-noted receiving, modifying and sending are illustratively repeated by the at least one processing device so as to ensure that each of the host devices of the host device cluster maintains a consistent stored version of the mapping for the given logical storage device.

(9) The cache entities of the storage system illustratively comprise respective components of the storage system that have respective local caches and/or respective portions of a global cache associated therewith. For example, the cache entities of the storage system may comprise respective storage controllers of the storage system with each such storage controller having a different local cache associated therewith.

(10) These and other illustrative embodiments include, without limitation, apparatus, systems, methods and computer program products comprising processor-readable storage media.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1 is a block diagram of an information processing system configured with functionality for synchronized mapping for cache-aware load balancing across multiple host devices in an illustrative embodiment.

(2) FIG. 2 is a flow diagram of a process for synchronized mapping for cache-aware load balancing across multiple host devices in an illustrative embodiment.

(3) FIG. 3 is a block diagram showing multiple layers of a layered system architecture that includes a multi-path layer with functionality for synchronized mapping for cache-aware load balancing across multiple host devices in an illustrative embodiment.

(4) FIG. 4 shows an example of storage-side performance information maintained by a storage system for use in cache-aware load balancing across multiple host devices in an illustrative embodiment.

(5) FIG. 5 is a block diagram of an information processing system configured with functionality for synchronized mapping for cache-aware load balancing across multiple host devices in another illustrative embodiment.

(6) FIGS. 6 and 7 are respective example signaling diagrams illustrating the operation of the FIG. 5 system in providing synchronized mapping across multiple host devices in an illustrative embodiment.

DETAILED DESCRIPTION

(7) Illustrative embodiments will be described herein with reference to exemplary information processing systems and associated computers, servers, storage devices and other processing devices. It is to be appreciated, however, that these and other embodiments are not restricted to the particular illustrative system and device configurations shown. Accordingly, the term “information processing system” as used herein is intended to be broadly construed, so as to encompass, for example, processing systems comprising cloud computing and storage systems, as well as other types of processing systems comprising various combinations of physical and virtual processing resources. An information processing system may therefore comprise, for example, at least one data center or other cloud-based system that includes one or more clouds hosting multiple tenants that share cloud resources, as well as other types of systems comprising a combination of cloud and edge infrastructure. Numerous different types of enterprise computing and storage systems are also encompassed by the term “information processing system” as that term is broadly used herein.

(8) FIG. 1 shows an information processing system **100** configured in accordance with an illustrative embodiment. The information processing system **100** comprises a computer system **101** that includes a plurality of host devices **102-1**, . . . **102-N**. The host devices **102** communicate over

a storage area network (SAN) **104** with at least one storage array **105**. The storage array **105** comprises a plurality of storage devices **106-1**, . . . **106-M** each storing data utilized by one or more applications running on one or more of the host devices **102**. The storage devices **106** are illustratively arranged in one or more storage pools. The storage array **105** and its associated storage devices **106** are an example of what is more generally referred to herein as a “storage system.” This storage system in the present embodiment is shared by the host devices **102**, and is therefore also referred to herein as a “shared storage system.” Other embodiments can include only a single host device, possibly configured to have exclusive use of the storage system.

(9) The host devices **102** illustratively comprise respective computers, servers or other types of processing devices capable of communicating with the storage array **105** of the SAN **104**. For example, at least a subset of the host devices **102** may be implemented as respective virtual machines of a compute services platform or other type of processing platform. The host devices **102** in such an arrangement illustratively provide compute services such as execution of one or more applications on behalf of each of one or more users associated with respective ones of the host devices **102**.

(10) The term “user” herein is intended to be broadly construed so as to encompass numerous arrangements of human, hardware, software or firmware entities, as well as combinations of such entities.

(11) Compute and/or storage services may be provided for users under a Platform-as-a-Service (PaaS) model, an Infrastructure-as-a-Service (IaaS) model, a Function-as-a-Service (FaaS) model and/or a Storage-as-a-Service (STaaS) model, although it is to be appreciated that numerous other cloud infrastructure arrangements could be used. Also, illustrative embodiments can be implemented outside of the cloud infrastructure context, as in the case of a stand-alone computing and storage system implemented within a given enterprise.

(12) The storage devices **106** of the storage array **105** of SAN **104** implement logical units (LUNs) configured to store objects for users associated with the host devices **102**. These objects can comprise files, blocks or other types of objects. The host devices **102** interact with the storage array **105** utilizing read and write commands as well as other types of commands that are transmitted over the SAN **104**. Such commands in some embodiments more particularly comprise Small Computer System Interface (SCSI) commands, although other types of commands can be used in other embodiments. A given IO operation as that term is broadly used herein illustratively comprises one or more such commands. References herein to terms such as “input-output” and “IO” should be understood to refer to input and/or output. Thus, an IO operation relates to at least one of input and output.

(13) Also, the term “storage device” as used herein is intended to be broadly construed, so as to encompass, for example, a logical storage device such as a LUN or other logical storage volume. A logical storage device can be defined in the storage array **105** to include different portions of one or more physical storage devices. Storage devices **106** may therefore be viewed as comprising respective LUNs or other logical storage volumes.

(14) Each of the host devices **102** illustratively has multiple paths to the storage array **105**, with at least one of the storage devices **106** of the storage array **105** being visible to that host device on a given one of the paths. A given one of the storage devices **106** may be accessible to the given host device over multiple paths.

(15) Different ones of the storage devices **106** of the storage array **105** illustratively exhibit different latencies in processing of IO operations. In some cases, the same storage device may exhibit different latencies for different ones of multiple paths over which that storage device can be accessed from a given one of the host devices **102**.

(16) The host devices **102**, SAN **104** and storage array **105** in the FIG. 1 embodiment are assumed to be implemented using at least one processing platform each comprising one or more processing devices each having a processor coupled to a memory. Such processing devices can illustratively

include particular arrangements of compute, storage and network resources. For example, processing devices in some embodiments are implemented at least in part utilizing virtual resources such as virtual machines (VMs) or Linux containers (LXC), or combinations of both as in an arrangement in which Docker containers or other types of LXCs are configured to run on VMs. (17) The host devices **102** and the storage array **105** may be implemented on respective distinct processing platforms, although numerous other arrangements are possible. For example, in some embodiments at least portions of the host devices **102** and the storage array **105** are implemented on the same processing platform. The storage array **105** can therefore be implemented at least in part within at least one processing platform that implements at least a subset of the host devices **102**.

(18) The SAN **104** may be implemented using multiple networks of different types to interconnect storage system components. For example, the SAN **104** may comprise a portion of a global computer network such as the Internet, although other types of networks can be part of the SAN **104**, including a wide area network (WAN), a local area network (LAN), a satellite network, a telephone or cable network, a cellular network, a wireless network such as a WiFi or WiMAX network, or various portions or combinations of these and other types of networks. The SAN **104** in some embodiments therefore comprises combinations of multiple different types of networks each comprising processing devices configured to communicate using Internet Protocol (IP) or other related communication protocols.

(19) As a more particular example, some embodiments may utilize one or more high-speed local networks in which associated processing devices communicate with one another utilizing Peripheral Component Interconnect express (PCIe) cards of those devices, and networking protocols such as InfiniBand (IB), Gigabit Ethernet or Fibre Channel (FC). Numerous alternative networking arrangements are possible in a given embodiment, as will be appreciated by those skilled in the art.

(20) The host devices **102** comprise respective sets of IO queues **110-1**, . . . **110-N** and respective MPIO drivers **112-1**, . . . **112-N**. The MPIO drivers **112** collectively comprise a multi-path layer of the host devices **102**. Path selection functionality for delivery of IO operations from the host devices **102** to the storage array **105** is provided in the multi-path layer by respective instances of path selection logic **114-1**, . . . **114-N** implemented within the MPIO drivers **112**. The multi-path layer further provides functionality for synchronized mapping across the host devices **102**. Such functionality is provided at least in part using respective instances of chunk mapping logic **115-1**, . . . **115-N** implemented within the MPIO drivers **112**.

(21) For example, each of the MPIO drivers **112** illustratively maintains, for each of one or more LUNs or other logical storage devices of the storage array **105**, a stored mapping of ranges of LBAs to cache entities of the storage array **105**, for use in cache-aware load balancing. The stored mappings are illustratively generated by the instances of chunk mapping logic **115** of the respective MPIO drivers **112**.

(22) An orchestrator device, also referred to herein as simply an “orchestrator,” interacts with the MPIO drivers **112** to ensure that the various stored mappings maintained by the different MPIO drivers **112** of the different host devices **102** are all consistent with one another, as will be described in more detail elsewhere herein. Such an arrangement is an example of what is more generally referred to herein as “synchronized mapping” for cache-aware load balancing across the multiple host devices **102**.

(23) The MPIO drivers **112** may comprise, for example, otherwise conventional MPIO drivers, such as PowerPath® drivers from Dell Technologies, suitably modified in the manner disclosed herein to provide functionality for synchronized mapping for cache-aware load balancing across the host devices **102**. Other types of MPIO drivers from other driver vendors may be suitably modified to incorporate functionality for synchronized mapping for cache-aware load balancing as disclosed herein.

(24) The host devices **102** can include additional or alternative components. For example, in some embodiments, the host devices **102** comprise respective local caches, implemented using respective memories of those host devices. A given such local cache can be implemented using one or more cache cards, possibly implementing caching techniques such as those disclosed in U.S. Pat. Nos. 9,201,803, 9,430,368 and 9,672,160, each entitled “System and Method for Caching Data,” and incorporated by reference herein. A wide variety of different caching techniques can be used in other embodiments, as will be appreciated by those skilled in the art. Other examples of memories of the respective host devices **102** that may be utilized to provide local caches include one or more memory cards or other memory devices, such as, for example, an NVMe over PCIe cache card, a local flash drive or other type of NVM storage drive, or combinations of these and other host memory devices.

(25) The system **100** further comprises an MPIO management station **116** that includes a processor **117** implementing interface logic **118**. The interface logic **118** is utilized to communicate with the host devices **102** and the storage array **105**. Such an MPIO management station **116** provides management functionality for the multi-path layer comprising the MPIO drivers **112** of the host devices **102**. In some embodiments, host device management software executing on the MPIO management station **116** interacts with storage array management software executing on the storage array **105**.

(26) The MPIO management station **116** may be viewed as one possible example of the above-noted orchestrator device that in illustrative embodiments is configured to synchronize stored mappings of logical block address ranges to respective cache entities of the storage array **105** across the host devices **102**, for each of one or more LUNs or other logical storage devices shared by the host devices **102**.

(27) The interface logic **118** in some embodiments further comprises mapping synchronization logic that is utilized by the MPIO management station **116** to synchronize the stored mappings across the host devices **102**, illustratively via interactions with their respective MPIO drivers as described elsewhere herein.

(28) The MPIO driver **112-1** is configured to deliver IO operations selected from its corresponding set of IO queues **110-1** to the storage array **105** via selected ones of multiple paths over the SAN **104**. The sources of the IO operations stored in the set of IO queues **110-1** illustratively include respective processes of one or more applications executing on the host device **102-1**. For example, IO operations can be generated by each of multiple processes of a database application running on the host device **102-1**. Such processes issue IO operations for delivery to the storage array **105** over the SAN **104**. Other types of sources of IO operations may be present in a given implementation of system **100**.

(29) A given IO operation is therefore illustratively generated by a process of an application running on the host device **102-1**, and is queued in a given one of the IO queues **110-1** of the host device **102-1** with other operations generated by other processes of that application, and possibly other processes of other applications.

(30) The paths from the host device **102-1** to the storage array **105** illustratively comprise paths associated with respective initiator-target pairs, with each initiator comprising a host bus adaptor (HBA) or other initiating entity of the host device **102-1** and each target comprising a port or other targeted entity corresponding to one or more of the storage devices **106** of the storage array **105**. As noted above, the storage devices **106** illustratively comprise LUNs or other types of logical storage devices.

(31) In some embodiments, the paths are associated with respective communication links between the host device **102-1** and the storage array **105** with each such communication link having a negotiated link speed. For example, in conjunction with registration of a given HBA to a switch of the SAN **104**, the HBA and the switch may negotiate a link speed. The actual link speed that can be achieved in practice in some cases is less than the negotiated link speed, which is a theoretical

maximum value.

(32) Various scheduling algorithms, load balancing algorithms and/or other types of algorithms can be utilized by the MPIO driver **112-1** in delivering IO operations from the IO queues **110-1** to the storage array **105** over particular paths via the SAN **104**. Each such IO operation is assumed to comprise one or more commands for instructing the storage array **105** to perform particular types of storage-related functions such as reading data from or writing data to particular logical volumes of the storage array **105**. Such commands are assumed to have various payload sizes associated therewith, and the payload associated with a given command is referred to herein as its “command payload.”

(33) A command directed by the host device **102-1** to the storage array **105** is considered an “outstanding” command until such time as its execution is completed in the viewpoint of the host device **102-1**, at which time it is considered a “completed” command. The commands illustratively comprise respective SCSI commands, although other command formats can be used in other embodiments. A given such command is illustratively defined by a corresponding command descriptor block (CDB) or similar format construct. The given command can have multiple blocks of payload associated therewith, such as a particular number of 512-byte SCSI blocks or other types of blocks.

(34) In illustrative embodiments to be described below, it is assumed without limitation that the initiators of a plurality of initiator-target pairs comprise respective HBAs of the host device **102-1** and that the targets of the plurality of initiator-target pairs comprise respective ports of the storage array **105**. Examples of such HBAs and storage array ports are illustrated in conjunction with the embodiment of FIG. 3.

(35) Selecting a particular one of multiple available paths for delivery of a selected one of the operations of the set of IO queues **110-1** is more generally referred to herein as “path selection.” Path selection as that term is broadly used herein can in some cases involve both selection of a particular IO operation and selection of one of multiple possible paths for accessing a corresponding logical device of the storage array **105**. The corresponding logical device illustratively comprises a LUN or other logical storage volume to which the particular IO operation is directed.

(36) It should be noted that paths may be added or deleted between the host devices **102** and the storage array **105** in the system **100**. For example, the addition of one or more new paths from host device **102-1** to the storage array **105** or the deletion of one or more existing paths from the host device **102-1** to the storage array **105** may result from respective addition or deletion of at least a portion of the storage devices **106** of the storage array **105**.

(37) Addition or deletion of paths can also occur as a result of zoning and masking changes or other types of storage system reconfigurations performed by a storage administrator or other user. Some embodiments are configured to send a predetermined command from the host device **102-1** to the storage array **105**, illustratively utilizing the MPIO driver **112-1**, to determine if zoning and masking information has been changed. The predetermined command can comprise, for example, a log sense command, a mode sense command, a “vendor unique command” or VU command, or combinations of multiple instances of these or other commands, in an otherwise standardized command format.

(38) In some embodiments, paths are added or deleted in conjunction with addition of a new storage array or deletion of an existing storage array from a storage system that includes multiple storage arrays, possibly in conjunction with configuration of the storage system for at least one of a migration operation and a replication operation.

(39) For example, a storage system may include first and second storage arrays, with data being migrated from the first storage array to the second storage array prior to removing the first storage array from the storage system.

(40) As another example, a storage system may include a production storage array and a recovery

storage array, with data being replicated from the production storage array to the recovery storage array so as to be available for data recovery in the event of a failure involving the production storage array.

(41) In these and other situations, path discovery scans may be repeated as needed in order to discover the addition of new paths or the deletion of existing paths.

(42) A given path discovery scan can be performed utilizing known functionality of conventional MPIO drivers, such as PowerPath® drivers.

(43) The path discovery scan in some embodiments may be further configured to identify one or more new LUNs or other logical storage volumes associated with the one or more new paths identified in the path discovery scan. The path discovery scan may comprise, for example, one or more bus scans which are configured to discover the appearance of any new LUNs that have been added to the storage array **105** as well to discover the disappearance of any existing LUNs that have been deleted from the storage array **105**.

(44) The MPIO driver **112-1** in some embodiments comprises a user-space portion and a kernel-space portion. The kernel-space portion of the MPIO driver **112-1** may be configured to detect one or more path changes of the type mentioned above, and to instruct the user-space portion of the MPIO driver **112-1** to run a path discovery scan responsive to the detected path changes. Other divisions of functionality between the user-space portion and the kernel-space portion of the MPIO driver **112-1** are possible. The user-space portion of the MPIO driver **112-1** is illustratively associated with an Operating System (OS) kernel of the host device **102-1**.

(45) For each of one or more new paths identified in the path discovery scan, the host device **102-1** may be configured to execute a host registration operation for that path. The host registration operation for a given new path illustratively provides notification to the storage array **105** that the host device **102-1** has discovered the new path.

(46) In the FIG. 1 embodiment, the storage array **105** comprises a plurality of storage controllers **120**, one or more sets of local caches **121**, and one or more instances of storage-side performance information **122**.

(47) The storage controllers **120** illustratively control the processing of IO operations received in the storage array **105** from the host devices **102**. For example, the storage controllers **120** illustratively manage the processing of read and write commands directed by the MPIO drivers **112** to particular ones of the storage devices **106**. The storage controllers **120** can be implemented as respective storage processors, storage directors or other storage system components configured to control storage system operations relating to processing of IO operations.

(48) In some embodiments, each of the storage controllers **120** has a different one of the local caches **121** associated therewith, although numerous alternative arrangements are possible. The storage controllers **120** are examples of what are more generally referred to herein as “cache entities” of the storage array **105**. A given “cache entity” of a storage system as the term is broadly used herein is intended to encompass a storage controller, storage processor, storage director or other system component that has a local cache or other type of cache allocated thereto or otherwise associated therewith. The caches illustratively include local caches **121**, although other types of caches can be used in other embodiments. These are considered examples of “storage caches” as they are implemented in the storage array **105** or other type of storage system, rather than in host devices **102**. The term “storage cache” as broadly used herein therefore refers to a cache implemented in a storage system.

(49) Accordingly, in some embodiments, the cache entities of the storage array **105** comprise respective storage controllers **120** of the storage array **105** that have respective local caches **121** associated therewith. As another example, in some embodiments, each of the storage controllers **120** has a different allocated portion of a global cache associated therewith, and such allocated portions of a global cache can be considered respective instances of local caches **121** associated with respective storage controllers **120**.

(50) The storage-side performance information **122** is illustratively maintained by the storage array **105** for respective LUNs or other logical storage devices, and can be arranged, for example, in the manner illustrated in the example of FIG. 4, or utilizing other types and arrangements of tables or other data structures.

(51) It is assumed that the MPIO drivers **112** of the respective host devices **102** can read or otherwise obtain at least portions of the storage-side performance information **122** directly or indirectly from the storage array **105**.

(52) For example, at least portions of the storage-side performance information **122** can be obtained by the host device **102-1** from the storage array **105** utilizing an in-band communication mechanism in which one or more predetermined commands in a designated storage protocol are sent from the host device **102-1** to the storage array **105**. Such predetermined commands can comprise, for example, a log sense command, a mode sense command, a “vendor unique command” or VU command, or combinations of multiple instances of these or other commands, in an otherwise standardized command format, such as a SCSI format, NVMe format or other type of format.

(53) A wide variety of different arrangements of commands may be used, as well as numerous associated timing techniques for repeatedly sending such commands from the host device **102-1** to the storage array **105**. A “command” as the term is broadly used herein can comprise a combination of multiple distinct commands, such as respective different commands for different ones of a plurality of logical storage devices. Separate commands are utilized in some embodiments to obtain storage-side performance information for respective logical storage devices. In other embodiments, a single command can be sent to obtain storage-side performance information for multiple logical storage devices. These and other references herein to sending a command from a host device to a storage array to obtain storage-side performance information are intended to be broadly construed.

(54) Additionally or alternatively, at least portions of the storage-side performance information **122** can be obtained by the host device **102-1** from the storage array **105** utilizing an out-of-band communication mechanism. For example, an out-of-band communication mechanism of this type can involve host management software of the host device **102-1** communicating with storage array management software of the storage array **105** over an IP network connection or other type of network connection. Such host management software can include software running on the MPIO management station **116**, in addition to or in place of software running on the individual host devices **102**.

(55) The MPIO driver **112-1** in some embodiments is configured to obtain at least portions of the storage-side performance information **122** maintained by storage array **105** in conjunction with processing of IO operations directed to the storage array **105** by host device **102-1** over SAN **104**, to dynamically select a particular one of a plurality of distinct load balancing policies available in the host device **102-1** based at least in part on the obtained storage-side performance information, and to apply the selected load balancing policy in directing additional IO operations from the host device **102-1** to the storage array **105**. In some embodiments, at least portions of the above-noted operations can be performed at least in part by other components of the host device **102-1** or by other processing devices, such as the MPIO management station **116**.

(56) The storage-side performance information is illustratively obtained from the storage array utilizing at least one of an in-band communication mechanism in which one or more commands in a designated storage protocol are sent from the host device **102-1** to the storage array **105**, and an out-of-band communication mechanism in which host management software of the host device **102-1** communicates with storage system management software of the storage array **105** over an IP connection or other type of network connection.

(57) At least one of the load balancing policies illustratively comprises a storage cache aware load balancing policy that causes different ones of the IO operations to be directed to different cache entities of the storage array **105** based at least in part on cache-related performance metrics of the

obtained storage-side performance information.

(58) The cache entities of the storage array **105** illustratively comprise respective components of the storage array **105** that have respective ones of the local caches **121** associated therewith. For example, the cache entities of the storage array **105** may comprise respective storage controllers **120** of the storage array **105** with each such storage controller having a different one of the local caches **121** associated therewith.

(59) In some embodiments, dynamically selecting a particular one of the plurality of distinct load balancing policies available in the host device based at least in part on the obtained storage-side performance information comprises automatically switching between at least a first load balancing policy that is storage cache aware and a second load balancing policy that is not storage cache aware.

(60) The first load balancing policy illustratively utilizes cache-related performance metrics of the obtained storage-side performance information in load balancing decisions, while the second load balancing policy illustratively does not utilize cache-related performance metrics of the obtained storage-side performance information in load balancing decisions. For example, the second load balancing policy can comprise an adaptive load balancing policy or another type of load balancing policy that is not storage cache aware.

(61) In some embodiments, the storage-side performance information comprises cache collision rates for respective cache entities of the storage array **105**.

(62) In such an arrangement, dynamically selecting a particular one of the plurality of distinct load balancing policies available in the host device based at least in part on the obtained storage-side performance information comprises automatically selecting a first load balancing policy that is storage cache aware to replace a second load balancing policy that is not storage cache aware responsive to cache collision rates for respective cache entities of the storage array **105** exceeding a specified threshold.

(63) In accordance with an example storage cache aware load balancing policy, the MPIO driver **112-1** in some embodiments is further configured to access for each of one or more of the storage devices **106** a stored mapping between ranges of LBAs of the storage device and respective ones of a plurality of cache entities of the storage array **105**. In some embodiments, each of the storage devices **106** has a corresponding distinct stored mapping of this type accessible to the MPIO driver **112-1**.

(64) For each of at least a subset of the IO operations, the MPIO driver **112-1** determines the particular storage device to which the IO operation is directed, and accesses the stored mapping for that storage device, also referred to as the “target storage device” of the IO operation. The MPIO driver **112-1** identifies a particular one of the cache entities, based at least in part on an LBA of the IO operation and the stored mapping, and selects a particular one of the paths for delivery of the IO operation to the storage array **105** based at least in part on the identified cache entity. Accordingly, the storage cache aware processing functionality of the MPIO driver **112-1** selects an appropriate path for delivery of the IO operation to the storage array **105**, based at least in part on a particular cache entity identified using an LBA of the IO operation and the stored mapping of the target storage device, and the IO operation is then delivered to the storage array **105** over the SAN **104** via the selected path.

(65) Additional examples of storage cache aware processing of IO operations are described below in conjunction with the embodiments of FIGS. 2 through 7.

(66) The above-noted stored mapping of LBA ranges to respective ones of the cache entities of the storage array, such as respective ones of the storage controllers **120**, is illustratively generated at least in part by the MPIO driver **112-1** and stored in one or more data structures of a kernel-space portion of an operating system of the host device **102-1**. Alternatively, other components of the host device **102-1** can be involved in at least one of the generating and the storage of the mapping for a given one of the storage devices **106**, using other types of data structures or storage arrangements.

(67) In some embodiments, a given such mapping is generated at least in part using information obtained from the storage array **105** and characterizing the cache entities of the storage array **105**, such as the storage controllers **120**.

(68) For example, the information characterizing the cache entities of the storage array **105** for use in generating the mapping illustratively comprises information identifying the storage controllers **120** of the storage array **105** that have respective local caches **121** associated therewith, information about the local caches **121**, and a chunk size associated with the local caches **121**, where the chunk size denotes a particular number of LBAs.

(69) The information characterizing the cache entities of the storage array **105** for use in generating the mapping can be obtained by the host device **102-1** from the storage array **105** utilizing an in-band communication mechanism in which one or more predetermined commands in a designated storage protocol are sent from the host device **102-1** to the storage array **105**. Such predetermined commands can comprise, for example, a log sense command, a mode sense command, a “vendor unique command” or VU command, or combinations of multiple instances of these or other commands, in an otherwise standardized command format, such as a SCSI format or other type of format.

(70) Additionally or alternatively, the information characterizing the cache entities of the storage array **105** for use in generating the mapping can be obtained by the host device **102-1** from the storage array **105** utilizing an out-of-band communication mechanism. For example, an out-of-band communication mechanism of this type can involve host management software of the host device **102-1** communicating with storage array management software of the storage array **105** over an IP network connection or other type of network connection.

(71) The MPIO driver **112-1** in some embodiments is configured to implement synchronized mapping in conjunction with cache-aware processing of IO operations utilizing storage cache aware load balancing as described above. In conjunction with the synchronized mapping functionality, the MPIO driver **112-1** separates LBAs of one or more storage devices of the storage array **105** into a plurality of ranges of LBAs using a designated chunk size, illustratively a predetermined initial chunk size, with the chunk size denoting a particular number of LBAs. A mapping of the type described previously is generated by assigning different ones of the ranges of LBAs to different ones of a plurality of cache entities of the storage array **105**. One or more of the cache entities of the storage array **105** may each be assigned multiple distinct ones of the ranges of LBAs.

(72) The MPIO driver **112-1** via its path selection logic **114-1** selects paths for delivery of respective IO operations from the host device **102-1** to the storage array **105** based at least in part on the assigning. The MPIO driver **112-1** via its chunk mapping logic **115-1** generates and maintains the stored mapping of LBA ranges to cache entities of the storage array **105**, illustratively with input from an orchestrator device such as the MPIO management station **116**, which in some embodiments is configured to support the synchronization of such stored mappings across the host devices **102**, through interaction with their respective MPIO drivers **112** as disclosed herein.

(73) In some embodiments, the MPIO management station **116** via its interface logic **118** is illustratively configured to store a mapping for each of a plurality of logical storage devices of the storage array **105**, with the mapping for a given one of the logical storage devices assigning different ranges of LBAs of the given logical storage device to respective different cache entities of the storage array **105**. The MPIO management station **116** receives at least one mapping-related communication from at least one of the host devices **102** that have shared access to the given logical storage device, modifies the stored mapping for the given logical storage device based at least in part on the at least one received mapping-related communication, and sends at least one mapping-related communication to at least one of the host devices **102** based at least in part on the modified stored mapping.

(74) The term “mapping-related communication” as used herein is intended to be broadly

construed, so as to encompass, for example, various types of notifications that may be sent from the host devices **102** to the MPIO management station **116**, as well as numerous other types of signals that relate in some way to mapping of logical block address ranges to cache entities of the storage array **105**.

(75) It is assumed in some embodiments that each of the host devices **102** has shared access to the given logical storage device, which may be a LUN or other type of logical storage volume. Other arrangements in which only a subset of the host devices **102** have such shared access to the given logical storage device are also possible.

(76) As indicated previously, the MPIO management station **116** is an example of an orchestrator device that is separate from the host devices **102**. The MPIO management station **116** is configured to interact with the MPIO drivers **112** of the respective host devices **102** in implementing the mapping synchronization functionality disclosed herein. Other types of orchestrator devices can be used in other embodiments. For example, an orchestrator device can be implemented in the form of one or more external servers that are configured to communicate with the host devices **102**.

(77) The above-noted receiving and sending of mapping-related communications between the host devices and the MPIO management station **116**, and the modifying of the stored mapping for the given logical storage device, are illustratively repeated on a continuous or periodic basis in order to ensure that each of the host devices **102** maintains a consistent stored version of the mapping for the given logical storage device.

(78) For example, in some embodiments, each of host devices **102** via its instance of chunk mapping logic **115** separately stores a version of the mapping for the given logical storage device and accesses its corresponding stored version of the mapping in conjunction with selecting paths for delivery of IO operations to the storage array **105**. Such mapping versions can be stored in persistent storage of respective ones of the host devices **102**, so as to be accessible to the respective MPIO driver **112** for use in path selection utilizing their corresponding instances of path selection logic **114**.

(79) The synchronized mapping functionality disclosed herein in some embodiments ensures that the various versions of the mapping stored by the host devices **102** and the MPIO management station **116** remain consistent with one another, illustratively through the actions of the MPIO management station **116** in operating as a conflict-eliminating intermediary between the host devices **102**.

(80) In some embodiments, receiving at least one mapping-related communication from at least one of the host devices **102** that have shared access to the given logical storage device comprises, for example, receiving a notification from the first host device **102-1** indicating that the first host device **102-1** has generated a new mapping entry.

(81) The MPIO management station **116** is illustratively configured to modify its stored mapping to include the new mapping entry responsive to a determination that there is no inconsistency between the new mapping entry and any corresponding existing mapping entry of the stored mapping.

(82) For example, the MPIO management station **116** in some embodiments is further configured to determine whether or not the new mapping entry is consistent with a corresponding existing mapping entry of the stored mapping.

(83) Responsive to an affirmative determination, the MPIO management station **116** sends the mapping-related communication to the first host device **102-1** with an instruction for the first host device **102-1** to utilize the new mapping entry in its version of the mapping.

(84) Responsive to a negative determination, the MPIO management station **116** sends the mapping-related communication to the first host device **102-1** with an instruction for the first host device **102-1** to utilize the corresponding existing mapping entry instead of the new mapping entry. Numerous other types and arrangements of mapping-related communications can be exchanged between the MPIO management station **116** and the host devices **102** in implementing mapping synchronization functionality as disclosed herein.

- (85) As another example, receiving at least one mapping-related communication from at least one of the host devices **102** that have shared access to the given logical storage device illustratively comprises receiving a notification from the first host device **102-1** indicating that all paths from the first host device **102-1** to a particular cache entity have failed.
- (86) The MPIO management station **116** is illustratively configured to determine whether or not at least a threshold number of the host devices **102** have each reported that all paths from that host device to the particular cache entity have failed.
- (87) Responsive to an affirmative determination, the MPIO management station **116** sends the mapping-related communication to each of the first host device **102-1** and one or more other ones of the host devices **102** with an instruction for that host device to clear its mapping entry for the particular cache entity and to create a new mapping entry for the particular cache entity.
- (88) Responsive to a negative determination, the MPIO management station **116** sends the mapping-related communication to at least the first host device **102-1** with an instruction for the first host device **102-1** to perform one or more designated tests, illustratively in order to confirm that the failure is not a transient failure and/or to identify the particular type of failure currently being encountered by the first host device **102-1** with respect to the particular cache entity, as an insufficient number of other ones of the host devices **102** are experiencing the same failure reported by the first host device **102-1**. The threshold number in some embodiments can include the first host device **102-1** and a single additional one of the host devices **102**, although other thresholds can be used in other embodiments.
- (89) Also, the particular mapping-related communications described above are examples only, and numerous other types of communications can be used to maintain consistency between the mappings of the host devices **102** and the MPIO management station **116** in the manner disclosed herein.
- (90) An example of a process illustrating one possible implementation of the above-noted synchronized mapping functionality will be described below in conjunction with the flow diagram of FIG. 2.
- (91) Again, in some embodiments, at least portions of the above-described synchronized mapping functionality can be performed at least in part by other components of the host device **102-1** or by other processing devices, such as the MPIO management station **116**, rather than by the MPIO driver **112-1**. References herein to such functionality being performed by “at least one processing device” can refer to the host device **102-1** or a portion thereof, or multiple processing devices possibly including multiple host devices **102** and/or the MPIO management station **116**. Numerous other arrangements of one or more processing devices can be used to implement the techniques disclosed herein.
- (92) As indicated previously, a mapping of the assignments of ranges of LBAs to respective ones of the cache entities is illustratively stored in one or more data structures of the host device **102-1** and accessed by the MPIO driver **112-1** in conjunction with selecting the paths for delivery of the IO operations to the storage array **105**.
- (93) Selecting paths for delivery of respective IO operations in some embodiments more particularly comprises accessing for a given one of the one or more storage devices **106** the above-noted stored mapping between ranges of LBAs of the storage device and respective ones of the cache entities of the storage array **105**. The path selection further involves, for each of at least a subset of the IO operations, identifying a particular one of the cache entities based at least in part on an LBA of the IO operation and the stored mapping, and selecting a particular one of a plurality of paths for delivery of the IO operation to the storage array **105** based at least in part on the identified cache entity.
- (94) In some embodiments, each of a plurality of LUNs or other logical storage devices of storage array **105** have their respective ranges of LBAs virtually divided or otherwise separated into repetitive chunks of chunk size B, where B is illustratively 4096 LBAs although other values could

be used. Each such chunk is mapped to a corresponding storage controller, storage processor, storage director or other cache entity of the storage array **105** associated with a particular one of the local caches of the storage array **105**.

(95) It should be noted that the chunk size can be specified in number of LBAs, as in the above example, or in a corresponding memory size associated with those LBAs, which is given by the number of LBAs times the memory size per LBA. Other types of chunk sizes can be used in other embodiments.

(96) In many situations, there could be certain IO operations whose starting LBA and ending LBA do not fall on the same chunk. Such IO operations are referred to herein as “overlapping” IOs since they each overlap two or more different chunks.

(97) As a result, regardless of which storage controller, storage processor, storage director or other cache entity of the storage array **105** is chosen as part of the path selection, there is an increased likelihood of performance penalties such as cache misses, as well as computation-intensive coordination activities between multiple storage controllers or directors, as would typically arise in the case of so-called “slot collisions.”

(98) Some embodiments herein configure host devices **102** to include functionality for dynamic chunk size adjustment based at least in part on detection of overlapping IO operations, utilizing the techniques disclosed in U.S. patent application Ser. No. 17/503,562, filed Oct. 18, 2021 and entitled “Dynamic Chunk Size Adjustment for Cache-Aware Load Balancing,” which is incorporated by reference herein.

(99) For example, illustrative embodiments disclosed in the above-cited U.S. patent application Ser. No. 17/503,562 configure the MPIO driver **112-1** to detect such overlapping IOs and to dynamically adjust the chunk size if the number of overlapping IOs detected within a particular time period exceeds a user-configurable threshold, which may be set to a particular percentage (e.g., 10% or 20%) of IOs within a given monitoring interval. In such an embodiment, if at least the threshold percentage of IOs processed by the MPIO driver **112-1** within the monitoring interval are overlapping IO operations, the MPIO driver **112-1** can automatically trigger an increase in the chunk size, such as a doubling in chunk size from 4096 LBAs to 8192 LBAs, although other types of dynamic increases in the chunk size can be used in other embodiments.

(100) This dynamic adjustment in chunk size based on detection of overlapping IO operations can potentially reduce the number of overlapping IO operations by about half, resulting in considerably reduced performance impacts from cache misses and slot collisions. Accordingly, IO latency is reduced and overall system performance is improved.

(101) Other illustrative embodiments disclosed herein utilize a fixed chunk size, rather than a dynamic chunk size. Accordingly, dynamic chunk size adjustment is not required in illustrative embodiments herein.

(102) In embodiments in which a given LUN or other storage device is shared by multiple ones of the host devices **102**, it is generally desirable to maintain consistency of the chunk size and other storage cache aware load balancing policy settings across all the host devices **102** that share that LUN or other storage device. This can be achieved using the MPIO management station **116**, which illustratively implements host management software for coordinating the operation of multi-pathing functionality across the host devices **102**.

(103) For example, the MPIO management station **116**, which illustratively has knowledge of shared LUNs and any changes in chunk size or other policy settings, can be responsible for applying such changes to all host devices **102** that share those particular LUNs. This functionality can involve use of an existing PowerPath® MPIO event alert mechanism or other suitable notifications from the MPIO driver **112-1** to inform the MPIO management station **116** of changes in policy parameter settings such as the change in chunk size for a given LUN. Such an arrangement allows the MPIO management station **116** to take proactive action to apply this change to all other host devices **102** sharing this LUN.

(104) The mapping of LBA ranges to cache entities of the storage array **105** is based on the chunk size, as described previously.

(105) In some embodiments, the mapping is generated for a given one of the storage devices **106** using the following steps: 1. Determining a size of the storage device in terms of a total number of LBAs within that storage device. 2. Separating the total number of LBAs of the storage device into a plurality of ranges of LBAs. 3. Assigning different ones of the ranges of LBAs to different ones of the cache entities of the storage array **105**.

(106) As noted above, the cache entities of the storage array **105** comprise respective ones of the storage controllers **120**. One or more of the cache entities are each assigned multiple distinct ones of the ranges of LBAs, utilizing the steps of the process described above. The resulting assignments of ranges of LBAs to respective ones of the cache entities comprise the mapping which is stored in the host device **102-1** so as to be accessible to the MPIO driver **112-1** in providing storage cache aware processing of IO operations as disclosed herein.

(107) Other mappings are similarly stored in each of the other host devices **102**, with coordination of the mappings being implemented via mapping synchronization functionality provided at least in part by MPIO management station **116** or other orchestrator device.

(108) In some embodiments, separating the total number of LBAs of the storage device into a plurality of ranges of LBAs in Step 2 above comprises separating the total number of LBAs into the plurality of ranges of LBAs using a designated chunk size, with the chunk size denoting a particular number of LBAs.

(109) As indicated previously, the chunk size may be dynamically adjustable based at least in part on detection of overlapping IO operations, illustratively utilizing the techniques disclosed in the above-cited U.S. patent application Ser. No. 17/503,562.

(110) A particular one of the cache entities is identified based at least in part on an LBA of the IO operation and the stored mapping using the following steps: 1. Determining an initial LBA and a transfer length for the IO operation. 2. Computing a target LBA based at least in part on the initial LBA and the transfer length. 3. Identifying a particular LBA range that includes the target LBA. 4. Utilizing the stored mapping to identify the particular one of the cache entities corresponding to the particular LBA range.

(111) The target LBA referred to in Steps 2 and 3 above is different than the initial LBA, and illustratively comprises an approximate midpoint LBA between the initial LBA and a final LBA as indicated by the transfer length. Other types of target LBAs can be used in identifying LBA ranges and corresponding cache entities.

(112) In some embodiments, the MPIO driver **112-1** is further configured to maintain multiple path sets associated with respective ones of the cache entities of the storage array **105**. In embodiments of this type, selecting a particular one of the paths for delivery of the IO operation to the storage array **105** in the path selection logic **114-1** based at least in part on the identified cache entity more particularly comprises identifying from the path sets at least one path set associated with the identified cache entity, and selecting the particular path from the identified path set.

(113) These and other functions related to storage cache aware processing of IO operations that are referred to herein as being performed by or under the control of the MPIO driver **112-1** can in some embodiments be performed at least in part outside of the MPIO driver **112-1**, such as in the host device **102-1** generally.

(114) References herein to an MPIO driver such as MPIO driver **112-1** being configured to access “stored mappings” are intended to be broadly construed, and should not be viewed as being limited to any particular storage arrangement, such as storage in a particular host device component or using a particular format. Accordingly, such storage can be in a memory or other storage device that is outside of the MPIO driver.

(115) Various conditions can lead to generation of new stored mappings, updating of existing stored mappings and/or deletion of existing stored mappings in illustrative embodiments. For example,

the addition of a new storage device to the set of storage devices **106** can trigger the creation of a new stored mapping for that storage device. Similarly, the deletion of an existing storage device from the set of storage devices **106** can trigger the removal of an existing stored mapping for that storage device from the system.

(116) It is to be appreciated that the term “stored mapping” as used herein is intended to be broadly construed, as so to encompass a wide variety of different types of stored information and associated storage arrangements for relating LBA ranges to cache entities of a storage array or other storage system.

(117) As indicated previously, absent use of the techniques for storage cache aware processing of IO operations in a multi-path layer as disclosed herein, performance of the storage system in processing IO requests received from a host device can be adversely impacted by cache-related issues of the storage system. For example, a host-based load balancing algorithm that is unaware of the particular cache entity configuration of a storage system can inadvertently cause an increased number of cache misses, thereby increasing the processing overhead of the storage system and degrading its performance.

(118) As an illustration of a problematic situation of the type described above, consider a simple example of a storage array having two controllers C1 and C2 with respective local caches CA1 and CA2. The two controllers C1 and C2 have respective ports P1 and P2. Assume that a write IO having an LBA offset of 100 is delivered from the multi-path layer of a host device to the storage array via port P1 of controller C1 and is therefore written to the local cache CA1 of controller C1. Further assume that the next write IO has the same LBA offset of 100. However, due to load balancing performed in conjunction with path selection by the multi-path layer of the host device, that write IO is delivered to the storage array via port P2 of controller C2 and is therefore written to the local cache CA2 of controller C2. The load balancing has therefore turned what might have been a cache hit in local cache CA1 into a cache miss in the local cache CA2. In addition, the previously-written data at LBA offset **100** in local cache CA1 will have to be marked as “dirty” and subsequently updated. Read IOs can cause similar adverse performance impacts within the storage array.

(119) Such drawbacks are advantageously overcome in illustrative embodiments herein by utilization of a multi-path layer comprising one or more of the MPIO drivers **112** to provide, for example, storage cache aware processing of IO operations as described above. These embodiments therefore provide a significant advance over conventional techniques that do not take storage cache arrangements of the storage system into account in path selection. For example, storage cache aware arrangements of the type disclosed herein can substantially reduce cache misses by better matching a load balancing algorithm or other type of path selection algorithm of a multi-path layer of a host device to the cache entity configuration of the storage system, leading to improved overall performance.

(120) The storage-side performance information **122** is not restricted to cache-related performance metrics, and can additionally or alternatively include a wide variety of other types of performance metrics, such as, for example, response time metrics, read/write ratios, write count, write pending count, bandwidth consumption and numerous others, illustrative maintained on a per-device basis, where the “device” refers to a LUN or other type of logical storage device. The above-described functionality associated with synchronized mapping is illustratively performed primarily by the MPIO driver **112-1** of the host device **102-1**, utilizing its chunk mapping logic **115-1** in cooperation with its path selection logic **114-1**, although other arrangements are possible. For example, these and other functions referred to herein as being performed by an MPIO driver can in other embodiments be performed at least in part by other host device components.

(121) An example of a set of storage-side performance information **122** maintained by the storage array **105** for respective LUNs or other logical storage devices is illustrated in FIG. 4. It is assumed that the storage array **105** maintains similar storage-side performance information for each of at

least a subset of the logical storage devices of the storage array **105**.

(122) As noted above, in some embodiments, storage-side performance information is obtained directly by the host device **102-1** from storage array **105**. Such an arrangement illustratively utilizes an in-band communication mechanism in which one or more commands in a designated storage protocol are sent from the host device **102-1** to the storage array **105**.

(123) It is possible in other embodiments that storage-side performance information can be obtained by the host device **102-1** indirectly from the storage array **105** utilizing an out-of-band communication mechanism via at least one intermediary device. An intermediary device in such an embodiment illustratively comprises the MPIO management station **116** of system **100**. The MPIO management station **116** can comprise a server configured to obtain the storage-side performance information directly from the storage array **105**. The MPIO management station **116** is illustratively connected directly to the storage array **105** using protocols such as SCSI, Internet SCSI (iSCSI), FC, NVMeoFC and/or others.

(124) The MPIO driver **112-1** on the host device **102-1** illustratively has connectivity to the MPIO management station **116**. The MPIO management station **116** in some embodiments implements PowerPath® Management Appliance (PPMA) functionality to obtain access to the storage array **105**. The MPIO driver **112-1** can obtain from the MPIO management station **116** certain types of storage array related information for use in various operations performed at least in part by the MPIO driver **112-1**, in addition to or in place of obtaining such information directly from the storage array **105**. Host multi-pathing software can be used to implement a multi-path layer comprising MPIO drivers **112** of respective host devices **102** as well as related management appliance software such as the above-noted PPMA of MPIO management station **116**. Such host multi-pathing software can be configured to obtain storage-side performance information or other types of information to facilitate synchronized mapping across multiple host devices as disclosed herein.

(125) Additionally or alternatively, PPMA functionality of the MPIO management station **116** can be used to ensure that all of the host devices **102** that share a particular LUN or other storage device all utilize the same chunk size and other policy parameters for storage cache aware processing of 10 operations targeting that storage device.

(126) Absent use of the mapping synchronization techniques disclosed herein, the host devices **102** that share access to LUNs or other logical storage devices of the storage array **105** may have inconsistent mappings, in which the same logical chunks of a given LUN may be mapped to different ones of the storage controllers **120** across the host devices **102**, which are illustratively arranged in a host device cluster. This can negatively impact performance in the system. Such negative impacts will generally be more visible in systems having large cluster configurations with large numbers of host devices sharing the storage array **105**. For example, a given host device which encounters path failures to any one of the storage array target ports may trigger a change in its logical chunk to cache entity mapping, potentially giving rise to inconsistencies between the mappings on the host devices **102**.

(127) In some embodiments, techniques are provided to handle situations in which one or more of the host devices **102** lose connectivity to at least one of the storage controllers **120**. Also, mechanisms are provided to monitor, correct and propagate storage controller mappings across the host devices **102**.

(128) Such functionality is provided in some embodiments utilizing the above-noted PPMA, which may be viewed as an example of what is more generally referred to herein as an “orchestrator.” For example, an otherwise conventional PPMA can be modified to provide mapping synchronization functionality as disclosed herein. In some embodiments, the MPIO management station **116** comprises such an orchestrator, also referred to herein as an “orchestrator device.”

(129) An orchestrator in some embodiments is configured to store and share the storage controller mapping information across the host devices **102**, while also performing other types of

management functionality for the host devices **102**. For example, the orchestrator in some embodiments is configured to collect and store the mapping information from the host devices **102**, and to share the mapping information across the host devices **102** so as to facilitate synchronization of such information between the host devices **102**. In some embodiments, this illustratively involves actively informing one or more of the host devices **102** to change the corresponding mapping in the case of a failure of a cache entity, so as to prevent workload imbalance across the storage controllers **120**.

(130) The orchestrator in some embodiments is configured to save this information in an associated database, providing for a consistent view across all the host devices **102**. Whenever there is a new mapping created or change in existing mapping on a given one of the host devices **102**, the MPIO driver **112** executing on the corresponding host device will notify the orchestrator through an event.

(131) For example, one or more host devices **102** may report to the orchestrator that all paths to a particular one of the storage controllers **120** are failed. The orchestrator can be configured to proactively notify all the host devices **102** of the host device cluster to initiate a change in their existing mapping based on at least a threshold number of host devices **102** reporting such failures. Additionally or alternatively, the orchestrator can be configured to wait for a predetermined timeout period before taking such actions, and/or may initiate tasks, possibly to be performed by one or more of the host devices **102**, to confirm that the issue is not transient.

(132) In some embodiments, a given one of the storage controllers **120** is considered an example cache entity of the storage array **105**, although other types of cache entities can be used. Multiple ones of the host devices **102** can each have multiple paths which are connected to a shared LUN from host initiators such as HBAs to target ports of the storage array **105**. The host devices **102** each have a connection to an orchestrator, possibly implemented at least in part through modification of the above-noted PPMA in the MPIO management station **116**, which can be used to provide communication between the host devices **102**, so as to facilitate synchronization of mappings across the host devices **102**.

(133) Load balancing policies implemented by the multi-path layer comprising MPIO drivers **112** of host devices **102** utilizing the above-described stored mappings are generally configured to provide balanced distribution of application workload to the storage array **105**. In some embodiments, the MPIO drivers **112** are configured to measure different host-side performance metrics such as, for example, number of active read and write blocks, read and write response times, read and write percentages, application workload patterns, and numerous others. However, excessive reliance on such host-side performance metrics can negatively impact application performance, due at least in part to the additional CPU cycles required to make such performance measurements.

(134) The use of storage-side performance metrics as disclosed herein not only alleviates these negative performance impacts on application performance, it can provide enhanced performance by avoiding slot collisions and other types of negative impacts attributable to cache-related entities of the storage array **105**.

(135) The slot collisions illustratively include contentions that arise when multiple directors or other storage controllers **120** of the storage array **105** access the same cache line. A storage cache aware load balancing policy as disclosed herein is advantageously configured to avoid such slot collisions. For example, a storage cache aware load balancing policy when enabled can ensure that slot collisions are avoided by distributing IO operations using a mapping of cache lines to each of a plurality of directors or other storage controllers **120** of the storage array **105**. The use of such a policy can provide significant performance improvements, particularly under conditions otherwise associated with a high level of slot collisions.

(136) Illustrative embodiments herein provide a trigger based on storage-side performance metrics, possibly supplemented by host-side performance metrics, that controls when the storage cache aware load balancing policy is applied in the host device. This avoids problems that might

otherwise arise if the storage cache aware load balancing policy were used under all conditions. For example, the storage cache aware load balancing policy typically works effectively in optimizing application workload when there is a relatively balanced mix of read and write IOs, but for some application workload patterns having higher write to read ratios, the storage cache aware load balancing policy can be less effective than other load balancing policies available in the host device.

(137) In some embodiments, the storage array **105** collects performance metrics for all storage access and provides a holistic view of all workloads coming from multiple hosts. Examples of parameters that can be captured by the storage array **105** include host IOPS, read/write response time, read/write percentages, read/write randomness, director/controller load, slot collisions, etc. These metrics are illustratively captured at one or more granularities, such as storage group, storage controller, storage device, array port, etc. These and other storage-side performance metrics are examples of what is more generally referred to herein as “storage-side performance information” and are made available to the MPIO drivers **112** of the host devices **102** in order to provide significantly enhanced load balancing, particularly with regard to dynamic selection between a load balancing policy that is storage cache aware and another load balancing policy that is not storage cache aware.

(138) In some embodiments, an example implementation comprises the following components, although additional or alternative components could be used: 1. Storage array management software (e.g., Unisphere for PowerMax™) running on storage array **105** or a related device. 2. MPIO software (e.g., PowerPath®) running in a software stack of host devices **102**. 3. PowerPath® Management Appliance (PPMA) implemented in MPIO management station **116**.

(139) Unisphere has access to and publishes different storage performance metrics at various granularities of host groups, storage groups, storage controllers, storage devices, storage array ports etc. Third-party software like PowerPath® can be configured to obtain access to these performance metrics from Unisphere through a Representational State Transfer (REST) application programming interface (API). Some of these performance metrics include host IOPS, read/write RT, workload pattern on storage device, sequential write/read, controller metrics, and cache metrics such as read/write hit percentage, read/write randomness, read/write percentage, etc.

(140) PowerPath® multi-pathing software manages storage devices and implements load balancing policies to distribute IO using its knowledge of application IO workload and visibility of storage endpoints. PowerPath® also has a feature called Performance Metrics Insight (PMI) that provides host-side view of storage performance.

(141) PPMA manages hosts running PowerPath® software and also has visibility of storage connectivity to each of the PowerPath® hosts along with access to PowerPath® host PMI data. PPMA can be configured (one time activity) to access Unisphere through REST API to get the performance metrics for the storage managed by PowerPath® on a host. PPMA can provide a centralized view of workload and performance metrics both from the storage array and from PowerPath® running on one or more hosts to build a performance monitoring and decision making platform that takes both these inputs to make sure customers can get the best of PowerPath® load balancing capabilities.

(142) The storage array **105** in some embodiments collects various performance metrics (e.g., cache/slot collisions, director/controller busy, read/write RT, etc.) and provides a detailed view of all workloads generated by multiple hosts attached to it via an interface such as Unisphere. Multipathing software like PowerPath® implements various load balancing policies such as an adaptive policy and a cache aware policy, and since no policy is perfect/optimal for all types of workload/environment situations, illustrative embodiments herein build intelligence into multipathing software intuitive enough to fine tune load balancing parameters and/or switch between such load balancing policy based on performance metrics generated by the storage array **105**.

(143) An illustrative embodiment in which portions of the functionality are provided using PPMA, implemented by MPIO management station **116** and modified to operate in accordance with techniques disclosed herein, is as follows. 1. PPMA collects storage performance statistics from Unisphere (e.g., using REST API) at the granularity of storage device being accessed by PowerPath® and also provides a holistic view of incoming workloads from all hosts on storage endpoints. 2. PPMA collects storage performance statistics from PowerPath® for the managed storage devices and workload distribution across storage endpoints. 3. PPMA monitors and pushes the key performance parameters to PowerPath® for analysis and load balancing policy selection and/or tuning of load balancing parameters.

(144) The above-noted PPMA arrangement illustratively implemented by MPIO management station **116** is illustratively configured to provide, in cooperation with the MPIO drivers **112**, a performance monitoring and decision making platform that has access to both storage-side performance metrics from the storage array **105** and host-side performance metrics from the host devices **102**.

(145) Examples of dynamic load balancing policy control implemented in illustrative embodiments include the following.

(146) In a first example, if the slot collision rate as detected by the MPIO driver in the obtained storage-side performance information is greater than an upper threshold (e.g., 20% default, user configurable) then the MPIO driver will select and apply a storage cache aware load balancing policy. Similarly, if the slot collision rate as detected by the MPIO driver in the obtained storage-side performance information is less than a lower threshold (e.g., 5% default, user configurable) then the MPIO driver will fall back to an adaptive load balancing policy (or another fallback policy as configured).

(147) In a second example, if the load imbalance across directors/controllers (globally) as detected by the MPIO driver in the obtained storage-side performance information is greater than an upper threshold (e.g., 30% default, user configurable) and if the host-side load imbalance (locally) is greater than a local threshold (e.g., 5%, user configurable) then the MPIO driver would select and apply an adaptive load balancing policy (or another load balancing policy that is not storage cache aware, as configured). Similarly, if the load imbalance across directors/controllers (globally) as detected by the MPIO driver in the obtained storage-side performance information is below a lower threshold (e.g., 5% default, user configurable) then the MPIO driver would select and apply the storage cache aware load balancing policy.

(148) The foregoing are just illustrative examples, and numerous other types of dynamic switching between load balancing policies based at least in part on storage-side performance information, possibly supplemented by host-side performance information, can be used in other embodiments.

(149) In some embodiments, the MPIO drivers **112** of the host devices **102** are configured to obtain an application IO workload pattern by deriving it from storage array performance metrics like read/write response time, sequential read/write percentage, read/write randomness, read/write percentages, skew, etc.

(150) Illustrative embodiments herein provide MPIO software with access to storage-side performance metrics to build an efficient load balancing policy based on these inputs. Such arrangements avoid excessive use of host-side software measurement logic that consumes additional CPU cycles resulting in overall performance impact when compared to having access to storage array metrics.

(151) Some embodiments are implemented in the form of an algorithm performed by cooperative interaction between the MPIO driver **112-1** of a given one of the host devices **102-1** and the storage array **105**, utilizing path selection logic **114-1** and chunk mapping logic **115-1**. Such an algorithm is more particularly implemented by cooperative interaction of instances of path selection logic **114** and instances of chunk mapping logic **115** of MPIO drivers **112** of host devices **102**, and storage controllers **120** of storage array **105**, possibly with involvement of one or more other components

of system **100**, such as MPIO management station **116**, which can provide the above-noted PPMA functionality, and is configured to operate as an orchestrator device in some embodiments.

(152) Illustrative embodiments provide significant advantages over conventional practice. For example, these embodiments provide an efficient algorithm for synchronized mapping across a cluster of multiple hosts in storage cache aware load balancing, as well as dynamic selection between multiple available load balancing policies that leverages storage-side performance information maintained in a storage array.

(153) These and other embodiments avoid the negative host performance implications of excessive cache misses and slot collisions. As a result, IO processing performance is improved, and the system can more easily meet its performance goals.

(154) The portions of the example techniques described above as being performed by a given MPIO driver **112-1** on a corresponding host device **102-1** can be similarly performed by other MPIO drivers **112** on respective other host devices **102**. Such MPIO drivers **112** illustratively form a multi-path layer comprising multi-pathing software of the host devices. Other types of host drivers can be used in other embodiments. For example, in some embodiments, at least a portion of the disclosed synchronized mapping functionality is carried out using one or more iSCSI drivers, or other types of non-MPIO host drivers.

(155) Additional examples of synchronized mapping arrangements will be described elsewhere herein in conjunction with the embodiments of FIGS. **2** through **7**. Other types of synchronized mapping arrangements can be used in other embodiments.

(156) These and other functions related to synchronized mapping that are referred to herein as being performed by or under the control of the MPIO driver **112-1** can in some embodiments be performed at least in part outside of the MPIO driver **112-1**, such as in the host device **102-1** generally.

(157) The above-described functions associated with synchronized mapping in the MPIO driver **112-1** in some embodiments are carried out at least in part under the control of its chunk mapping logic **115-1**, illustratively operating in cooperation with path selection logic **114-1**. For example, the chunk mapping logic **115-1** is illustratively configured to control generation and maintenance of stored mappings for respective logical storage devices. In other embodiments, one or more such steps can be more generally performed by the host device **102-1**.

(158) It is assumed that each of the other MPIO drivers **112** is configured in a manner similar to that described above and elsewhere herein for the first MPIO driver **112-1**. The other host devices **102** of the system **100** are therefore also configured to communicate over the SAN **104** with the storage array **105**. The MPIO drivers **112** of such other host devices are each similarly configured to deliver IO operations from its corresponding one of the sets of IO queues **110** to the storage array **105** over selected paths through the SAN **104**, and to perform the disclosed functionality for synchronized mapping, illustratively through interaction with the MPIO management station **116**.

(159) Accordingly, functionality described above in the context of the first MPIO driver **112-1** and the first host device **102-1** is assumed to be similarly performed by each of the other MPIO drivers **112-2** through **112-N** and/or more generally by their respective host devices **102-2** through **102-N**.

(160) The MPIO drivers **112** may be otherwise configured utilizing well-known MPIO functionality. Such conventional MPIO functionality is suitably modified in illustrative embodiments disclosed herein to support synchronized mapping.

(161) Although in some embodiments certain commands used by the host devices **102** to communicate with the storage array **105** illustratively comprise SCSI commands, other types of commands and command formats can be used in other embodiments. For example, some embodiments can implement IO operations utilizing command features and functionality associated with NVMe Express (NVMe), as described in the NVMe Base Specification, Revision 2.0b, Jan. 6, 2022, which is incorporated by reference herein. Other storage protocols of this type that may be utilized in illustrative embodiments disclosed herein include NVMe over Fabric, also referred to as

NVMeoF.

(162) The storage array **105** in the present embodiment is assumed to comprise a persistent memory that is implemented using a flash memory or other type of non-volatile memory of the storage array **105**. More particular examples include NAND-based flash memory or other types of non-volatile memory such as resistive RAM, phase change memory, spin torque transfer magnetoresistive RAM (STT-MRAM) and Intel Optane™ devices based on 3D)(Point™ memory. The persistent memory is further assumed to be separate from the storage devices **106** of the storage array **105**, although in other embodiments the persistent memory may be implemented as a designated portion or portions of one or more of the storage devices **106**. For example, in some embodiments the storage devices **106** may comprise flash-based storage devices, as in embodiments involving all-flash storage arrays, or may be implemented in whole or in part using other types of non-volatile memory.

(163) As mentioned above, communications between the host devices **102** and the storage array **105** may utilize PCIe connections or other types of connections implemented over one or more networks. For example, illustrative embodiments can use interfaces such as Internet SCSI (iSCSI), Serial Attached SCSI (SAS) and Serial ATA (SATA). Numerous other interfaces and associated communication protocols can be used in other embodiments.

(164) The storage array **105** in some embodiments may be implemented as part of cloud infrastructure in the form of a cloud-based system such as an Amazon Web Services (AWS) system. Other examples of cloud-based systems that can be used to provide at least portions of the storage array **105** and possibly other portions of system **100** include Google Cloud Platform (GCP) and Microsoft Azure.

(165) The storage devices **106** of the storage arrays **105** can be implemented using solid state drives (SSDs). Such SSDs are implemented using non-volatile memory (NVM) devices such as flash memory. Other types of NVM devices that can be used to implement at least a portion of the storage devices **106** include non-volatile random access memory (NVRAM), phase-change RAM (PC-RAM) and magnetic RAM (MRAM). These and various combinations of multiple different types of NVM devices or other storage devices may also be used. For example, hard disk drives (HDDs) can be used in combination with or in place of SSDs or other types of NVM devices. Accordingly, numerous other types of electronic or magnetic media can be used in implementing at least a subset of the storage devices **106**.

(166) The storage array **105** may additionally or alternatively be configured to implement multiple distinct storage tiers of a multi-tier storage system. By way of example, a given multi-tier storage system may comprise a fast tier or performance tier implemented using flash storage devices or other types of SSDs, and a capacity tier implemented using HDDs, possibly with one or more such tiers being server based. A wide variety of other types of storage devices and multi-tier storage systems can be used in other embodiments, as will be apparent to those skilled in the art. The particular storage devices used in a given storage tier may be varied depending on the particular needs of a given embodiment, and multiple distinct storage device types may be used within a single storage tier. As indicated previously, the term “storage device” as used herein is intended to be broadly construed, and so may encompass, for example, SSDs, HDDs, flash drives, hybrid drives or other types of storage products and devices, or portions thereof, and illustratively include logical storage devices such as LUNs.

(167) As another example, the storage array **105** may be used to implement one or more storage nodes in a cluster storage system comprising a plurality of storage nodes interconnected by one or more networks.

(168) It should therefore be apparent that the term “storage array” as used herein is intended to be broadly construed, and may encompass multiple distinct instances of a commercially-available storage array. For example, the storage array **105** may comprise one or more storage arrays such as one or more PowerMax™ storage arrays, commercially available from Dell Technologies.

(169) Other types of storage products that can be used in implementing a given storage system in illustrative embodiments include software-defined storage, cloud storage, object-based storage and scale-out storage. Combinations of multiple ones of these and other storage types can also be used in implementing a given storage system in an illustrative embodiment.

(170) In some embodiments, a storage system comprises first and second storage arrays arranged in an active-active configuration. For example, such an arrangement can be used to ensure that data stored in one of the storage arrays is replicated to the other one of the storage arrays utilizing a synchronous replication process. Such data replication across the multiple storage arrays can be used to facilitate failure recovery in the system **100**. One of the storage arrays may therefore operate as a production storage array relative to the other storage array which operates as a backup or recovery storage array.

(171) It is to be appreciated, however, that embodiments disclosed herein are not limited to active-active configurations or any other particular storage system arrangements. Accordingly, illustrative embodiments herein can be configured using a wide variety of other arrangements, including, by way of example, active-passive arrangements, Asymmetric Logical Unit Access (ALUA) arrangements and/or Asymmetric Namespace Access (ANA) arrangements.

(172) These and other storage systems can be part of what is more generally referred to herein as a processing platform comprising one or more processing devices each comprising a processor coupled to a memory. A given such processing device may correspond to one or more virtual machines or other types of virtualization infrastructure such as Docker containers or other types of LXC's. As indicated above, communications between such elements of system **100** may take place over one or more networks.

(173) The term “processing platform” as used herein is intended to be broadly construed so as to encompass, by way of illustration and without limitation, multiple sets of processing devices and one or more associated storage systems that are configured to communicate over one or more networks. For example, distributed implementations of the host devices **102** are possible, in which certain ones of the host devices **102** reside in one data center in a first geographic location while other ones of the host devices **102** reside in one or more other data centers in one or more other geographic locations that are potentially remote from the first geographic location. Thus, it is possible in some implementations of the system **100** for different ones of the host devices **102** to reside in different data centers than the storage array **105**.

(174) Numerous other distributed implementations of the host devices **102** and/or the storage array **105** are possible. Accordingly, the storage array **105** can also be implemented in a distributed manner across multiple data centers.

(175) It is to be appreciated that these and other features of illustrative embodiments are presented by way of example only, and should not be construed as limiting in any way. Accordingly, different numbers, types and arrangements of system components such as host devices **102**, SAN **104**, storage array **105**, storage devices **106**, sets of IO queues **110**, and MPIO drivers **112**, including their corresponding instances of path selection logic **114** and chunk mapping logic **115**, can be used in other embodiments.

(176) It should also be understood that the particular sets of modules and other components implemented in the system **100** as illustrated in FIG. **1** are presented by way of example only. In other embodiments, only subsets of these components, or additional or alternative sets of components, may be used, and such components may exhibit alternative functionality and configurations.

(177) The operation of the information processing system **100** will now be described in further detail with reference to the flow diagram of the illustrative embodiment of FIG. **2**. The process as shown includes steps **200** through **208**, and is suitable for use in the system **100** but is more generally applicable to other types of systems comprising multiple host devices that share access to at least one logical storage device of a storage system. The storage system in this embodiment is

assumed to comprise at least one storage array having a plurality of storage devices. The storage devices can include logical storage devices such as LUNs or other logical storage volumes.

(178) The steps of the FIG. 2 process are illustratively performed at least in part by or under the control of an orchestrator device, such as, for example, the MPIO management station **116**, also referred to herein as an “orchestrator,” which interacts with multiple host devices of a host device cluster. For example, portions of the process may be carried out under the control of software, firmware and/or hardware of an orchestrator device interacting with the multiple host devices of the host device cluster.

(179) Accordingly, the FIG. 2 process may be viewed as an example of an algorithm performed in part by the MPIO management station **116** interacting with the instances of chunk mapping logic **115** of the MPIO drivers **112** of the host devices **102**. It is to be appreciated that other arrangements of host device components, storage array components and/or other system components can be configured to perform at least portions of one or more of the steps in other embodiments. For example, in some embodiments, the disclosed orchestrator device may be implemented in a particular one of the host devices **102**, or in a distributed manner across multiple ones of the host devices **102**.

(180) In step **200**, the orchestrator maintains a stored mapping of ranges of LBAs of a logical storage device to cache entities of a storage array. The mapping utilizes a designated chunk size, where the chunk size illustratively denotes a particular number of LBAs. The chunk size may denote a particular number of LBAs in a variety of different ways. For example, the chunk size can be specified in terms of a number of LBAs or in terms of an amount of memory represented by that number of LBAs. Other techniques can be used to specify the chunk size in other embodiments, and the term “chunk size” as used herein is therefore intended to be broadly construed.

(181) Such a mapping can be initially generated in each of one or more of the host devices **102**, and communicated to the orchestrator. Additionally or alternatively, such a mapping can be generated in a particular one of the host devices **102**, such as host device **102-1**, and then communicated to one or more other ones of the host devices **102**, possibly via the MPIO management station **116**. As another example, the mapping can initially be generated by the MPIO management station **116** and then communicated or otherwise made available to each of the host devices **102**.

(182) In some embodiments, it is desirable to use the same chunk size and mapping in each of the host devices **102**, and thus the mapping can be performed by one of the host devices **102** and then communicated to all of the other host devices **102**, possibly via the MPIO management station **116**, or can be performed by the MPIO management station **116** and then communicated or otherwise made available by the MPIO management station **116** to each of the host devices **102**. Numerous other mapping arrangements can be used in other embodiments, and the disclosed embodiments are not limited in this regard.

(183) In step **202**, the orchestrator receives one or more mapping-related communications from at least one host. For example, the MPIO management station **116** can receive communications from one or more of the host device **102** that propose various changes to the mapping of LBA ranges to cache entities for the given logical storage device.

(184) In step **204**, the orchestrator determines whether the received communications indicate a mapping inconsistency. Such a mapping inconsistency is illustratively determined relative to a stored mapping maintained by the orchestrator based on its interactions with multiple hosts of the host cluster.

(185) In step **206**, responsive to a mapping inconsistency being detected between two hosts, the process moves to step **208**, and otherwise returns to step **202**.

(186) In step **208**, the orchestrator sends one or more mapping-related communications to instruct one or more hosts in order to eliminate the detected inconsistency. The process then returns to step **200** as indicated.

(187) Multiple additional instances of the FIG. 2 process may be performed for respective

additional logical storage devices that are shared by the multiple hosts of the host cluster. As indicated above, the MPIO management station **116** can operate as the orchestrator in the FIG. 2 process, and in some embodiments is configured to ensure that all of the host devices **102** are utilizing the same chunk size and the same mapping of LBA ranges of the logical storage device to cache entities of the storage array **105**.

(188) The particular processing operations and other system functionality described in conjunction with the flow diagram of FIG. 2 are presented by way of illustrative example only, and should not be construed as limiting the scope of the disclosure in any way. Alternative embodiments can use other types of processing operations involving host devices, storage systems and functionality for synchronized mapping. For example, the ordering of the process steps may be varied in other embodiments, or certain steps may be performed at least in part concurrently with one another rather than serially. Also, one or more of the process steps may be repeated periodically, or multiple instances of the process can be performed in parallel with one another in order to implement a plurality of different synchronized mapping arrangements within a given information processing system.

(189) Functionality such as that described in conjunction with the flow diagram of FIG. 2 can be implemented at least in part in the form of one or more software programs stored in memory and executed by a processor of a processing device such as a computer or server. As will be described below, a memory or other storage device having executable program code of one or more software programs embodied therein is an example of what is more generally referred to herein as a “processor-readable storage medium.”

(190) Referring now to FIG. 3, another illustrative embodiment is shown. In this embodiment, an information processing system **300** comprises host-side elements that include application processes **311**, path selection logic **314** and chunk mapping logic **315**, and storage-side elements that include cache control logic **321** and performance monitor **322**. The cache control logic **321** is configured to manage local caches of one or more storage arrays, although such local caches are not explicitly shown in the figure. There may be separate instances of one or more such elements associated with each of a plurality of system components such as host devices and storage arrays of the system **300**.

(191) In a manner similar to that described elsewhere herein, storage-side performance metrics associated with local caches and additional or alternative components of one or more storage arrays are generated by the performance monitor **322**, and may be obtainable by one or more host devices via their respective MPIO drivers and utilized by the MPIO drivers to select a particular load balancing policy responsive to the storage-side performance metrics. The selected load balancing policy illustratively determines the manner in which IO operations are sent to particular logical storage devices of the one or more storage arrays over paths selected by the path selection logic **314**. In some embodiments herein, it is assumed that a storage cache aware load balancing policy with synchronized mapping is selected. Other types of load balancing policies, such as load balancing policies that are not storage cache aware, can be selected by the MPIO drivers, based at least in part on the storage-side performance metrics.

(192) The system **300** is configured in accordance with a layered system architecture that illustratively includes a host device processor layer **330**, an MPIO layer **332**, an HBA layer **334**, a switch fabric layer **336**, a storage array port layer **338** and a storage array processor layer **340**. The host device processor layer **330**, the MPIO layer **332** and the HBA layer **334** are associated with one or more host devices, the switch fabric layer **336** is associated with one or more SANs or other types of networks, and the storage array port layer **338** and storage array processor layer **340** are associated with one or more storage arrays (“SAs”).

(193) The system **300** in this embodiment implements synchronized mapping utilizing one or more MPIO drivers of the MPIO layer **332**, and associated instances of chunk mapping logic **315**. The application processes **311** generate IO operations that are processed by the MPIO layer **332** for delivery to the one or more storage arrays. Paths are determined by the path selection logic **314** for

sending such IO operations to the one or more storage arrays. These IO operations are sent to the one or more storage arrays in accordance with a load balancing policy that may be selected by the MPIO driver based at least in part on storage-side performance metrics obtained by the MPIO driver for respective logical storage devices. For example, as indicated above, a storage cache aware load balancing policy with synchronized mapping is illustratively selected. In accordance with such a storage cache aware load balancing policy, the chunk mapping logic **315** operates in cooperation with the path selection logic **314** in one or more MPIO drivers of the MPIO layer **332** to control delivery of IO operations in the manner disclosed herein. Again, a variety of other types of load balancing policies, such as load balancing policies that are not storage cache aware, can be selected by the MPIO drivers, based at least in part on the storage-side performance metrics.

(194) The MPIO layer **332** is an example of what is also referred to herein as a multi-path layer, and comprises one or more MPIO drivers implemented in respective host devices. Each such MPIO driver illustratively comprises respective instances of path selection logic **314** and chunk mapping logic **315** configured as previously described. Additional or alternative layers and logic arrangements can be used in other embodiments.

(195) For example, the MPIO layer **332** illustratively comprises at least one MPIO management station or other type of orchestrator that operates as described elsewhere herein to ensure that a consistent mapping is utilized by different host devices of a host device cluster, for a given logical storage device shared by those host devices, in implementing storage cache aware load balancing in system **300**.

(196) The cache control logic **321** implemented in the storage array processor layer **340** controls the configuration and operation of caches associated with respective cache entities of the system **300**, such as multiple local caches associated with respective ones of a plurality of storage controllers of the system **300**. For example, the cache control logic **321** can include functionality for allocating available cache resources among a plurality of cache entities of a given storage array of the system **300**. It is also possible in some embodiments that the cache control logic **321** can include multiple distinct cache control logic instances for multiple sets of cache entities of respective ones of a plurality of storage arrays of the system **300**.

(197) In the system **300**, path selection logic **314** is configured to select different paths for sending IO operations from a given host device to a storage array. These paths as illustrated in the figure include a first path from a particular HBA denoted HBA**1** through a particular switch fabric denoted SF**1** to a particular storage array port denoted PORT**1**, and a second path from another particular HBA denoted HBA**2** through another particular switch fabric denoted SF**2** to another particular storage array port denoted PORT**2**.

(198) These two particular paths are shown by way of illustrative example only, and in many practical implementations there will typically be a much larger number of paths between the one or more host devices and the one or more storage arrays, depending upon the specific system configuration and its deployed numbers of HBAs, switch fabrics and storage array ports. For example, each host device in the FIG. **3** embodiment can illustratively have a set of n paths to a shared storage array, or alternatively different ones of the host devices can have different numbers and types of paths to the storage array.

(199) The path selection logic **314** of the MPIO layer **332** in this embodiment selects paths for delivery of IO operations to the one or more storage arrays having the storage array ports of the storage array port layer **338**. More particularly, the path selection logic **314** determines appropriate paths over which to send particular IO operations to particular logical storage devices of the one or more storage arrays, based at least in part on a stored mapping that is determined utilizing a particular chunk size.

(200) It is assumed in this embodiment that the host devices through their respective MPIO drivers and respective instances of chunk mapping logic **315** provide functionality for synchronized mapping as disclosed herein, possibly with involvement of other host device components such as

the path selection logic **314**, and additionally or alternatively with involvement of one or more other system components such as an MPIO management station or other arrangement of one or more external servers.

(201) Some implementations of the system **300** can include a relatively large number of host devices (e.g., 1000 or more host devices), although as indicated previously different numbers of host devices, and possibly only a single host device, may be present in other embodiments. Each of the host devices is typically allocated with a sufficient number of HBAs to accommodate predicted performance needs. In some cases, the number of HBAs per host device is on the order of 4, 8 or 16 HBAs, although other numbers of HBAs could be allocated to each host device depending upon the predicted performance needs. A typical storage array may include on the order of 128 ports, although again other numbers can be used based on the particular needs of the implementation. The number of host devices per storage array port in some cases can be on the order of 10 host devices per port. The HBAs of the host devices are assumed to be zoned and masked to the storage array ports in accordance with the predicted performance needs, including user load predictions.

(202) A given host device of system **300** can be configured to initiate an automated path discovery process to discover new paths responsive to updated zoning and masking or other types of storage system reconfigurations performed by a storage administrator or other user. For certain types of host devices, such as host devices using particular operating systems such as Windows, ESX or Linux, automated path discovery via the MPIO drivers of a multi-path layer is typically supported. Other types of host devices using other operating systems such as AIX in some implementations do not necessarily support such automated path discovery, in which case alternative techniques can be used to discover paths.

(203) Additional illustrative embodiments will now be described. It is assumed in these embodiments that the MPIO driver of a given host device provides at least portions of the disclosed synchronized mapping functionality, possibly with involvement of one or more other host device components.

(204) Referring now to FIG. 4, an example of storage-side performance information **400** maintained by a storage system such as storage array **105** is shown. The storage-side performance information **400** is illustratively maintained by the storage array **105** for particular LUNs that are denoted LUN 1, LUN 2, . . . LUN X in the figure. These LUNs may comprise all of the LUNs of the storage array **105**, or a designated subset of the LUNs of the storage array **105**. Different sets of storage-side performance information similar to storage-side performance information **400** may be maintained for respective different sets of LUNs of the storage array **105**. The storage-side performance information **400** may be viewed as one possible example of storage-side performance information **122** of storage array **105** in the FIG. 1 embodiment, illustratively generated by a corresponding performance monitor not explicitly shown in FIG. 1. Similar storage-side performance information may be generated by performance monitor **322** in the FIG. 3 embodiment.

(205) The storage-side performance information **400** more particularly includes cache-related performance metrics for each of the LUNs denoted LUN 1 through LUN X. The storage-side performance measurements in some embodiments can additionally or alternatively comprise other types of performance metrics, such as response time (RT) measurements compiled by the storage array **105** over a designated time period for IO operations directed to the corresponding LUNs. Additional or alternative performance metrics that can be used in generating the performance measurements that are part of the storage-side performance information include, for example, at least one of a write count, a write pending count and a bandwidth consumption for IO operations directed to the corresponding LUNs.

(206) At least portions of the storage-side performance information **400** are obtainable by the MPIO drivers of respective host devices and is utilized by those MPIO drivers in selecting a particular type of load balancing policy for use in delivering IO operations to the storage array as disclosed herein. For example, a storage cache aware load balancing policy with dynamically

adjustable chunk size may be selected.

(207) The particular storage-side performance information arrangement shown in FIG. 4 is only an example, and numerous other types and arrangements of storage-side performance information can be maintained by a storage system in other embodiments. Terms such as “storage-side performance information” as used herein are intended to be broadly construed.

(208) As indicated above, the MPIO driver of a given one of the host devices **102** can obtain storage-side performance information from the storage array **105** by sending commands to the storage array **105**, in the manner previously described.

(209) The storage-side performance information **400** is illustratively stored in a storage array memory accessible to a multi-path layer of the host device. Other types of data structures can be used in other embodiments.

(210) On the host side, at least portions of such storage-side performance information obtained from the storage array and characterizing one or more logical storage devices are illustratively stored in a data structure of a kernel-space portion of the MPIO driver **112-1**, although other storage arrangements with other types of data structures can be used.

(211) In some embodiments, an MPIO driver obtains performance information directly from a storage array, through interaction with the storage array as described elsewhere herein. For example, MPIO driver **112-1** is illustratively configured to obtain the storage-side performance information directly from the storage array **105** utilizing an in-band communication mechanism in which one or more commands in a designated storage protocol are sent from the MPIO driver **112-1** or another host device component to the storage array **105**.

(212) However, it is possible in some embodiments that one or more MPIO drivers of one or more host devices can obtain the storage-side performance information other than through direct communication with the storage array. Such arrangements illustratively involve obtaining the connectivity information indirectly from the storage array utilizing an out-of-band communication mechanism via at least one intermediary device, such as the MPIO management station **116**, which may comprise one or more servers.

(213) Additional illustrative embodiments will now be described with reference to FIGS. 5, 6 and 7. It is assumed in these embodiments that an orchestrator device coordinates the mappings utilized by multiple host devices for implementing the disclosed storage cache aware processing functionality. The storage cache aware processing is an example of what is also referred to herein as a storage cache aware load balancing policy, illustratively selected for use based at least in part on obtained storage-side performance information, possibly supplemented in some embodiments by host-side performance information. It is to be appreciated that the particular features and functionality of the embodiments of FIGS. 5, 6 and 7 as described below can be varied in other embodiments.

(214) Referring initially to FIG. 5, an information processing system **500** comprises a host device cluster that includes three host devices **502-1**, **502-2** and **502-3**, also denoted Host Device **1**, Host Device **2** and Host Device **3**, respectively. A given host device cluster as that term is broadly used herein can include more or fewer host devices. Each of the host devices **502** is assumed to communicate with a storage array **505** over respective sets of paths in order to obtain access to a particular shared LUN **506-1**. It is to be appreciated that there may be numerous additional shared LUNs in the storage array **505**, although only the single shared LUN **506-1** is shown in the figure for clarity and simplicity of illustration. The paths over which the host devices **502** communicate with the storage array are associated with respective initiator-target pairs, as described elsewhere herein. The initiators illustratively comprise respective HBAs of the host devices **502**, and the targets illustratively comprise respective ports of the storage array **505**. The paths in this embodiment may be arranged in respective path sets. The host devices **502-1**, **502-2** and **502-3** comprise respective MPIO drivers **512-1**, **512-2** and **512-3**, each configured as described elsewhere herein to provide storage cache aware control of delivery of IO operations to the storage array **505**.

over selected ones of the paths of the path sets.

(215) The storage cache aware functionality of the MPIO drivers **512** is implemented at least in part utilizing, for the shared LUN **506-1**, a corresponding mapping of LBA ranges of the shared LUN **506-1** to particular cache entities of the storage array **505**.

(216) The system **500** further comprises an orchestrator **516** that includes mapping synchronization logic **518**. The mapping synchronization logic **518** is configured to control communications between the orchestrator **516** and the MPIO drivers **512** of the respective host devices **502** in order to ensure that a consistent mapping is used by each of the host devices **502** in accessing the shared LUN **506-1**.

(217) The cache entities of the storage array **505** in this embodiment include first and second storage controllers **520-1** and **520-2**, also denoted as Controller **1** and Controller **2**, respectively, which are associated with respective first and second caches **521-1** and **521-2**, also denoted as Cache **1** and Cache **2**, respectively. The first and second caches **521-1** and **521-2** may comprise separate local caches of the respective first and second storage controllers **520-1** and **520-2**, and/or different allocated portions of a global cache of the storage array **505**. The storage controllers **520** are examples of what are more generally referred to herein as respective “cache entities” of the storage array **505**, and may in some cases include their respective caches **521-1** and **521-2**.

Alternatively, the caches **521-1** and **521-2** are arranged outside of their respective storage controllers, as in the particular arrangement illustrated in the figure. The storage controllers **520** in some embodiments are implemented as storage processors, storage directors, or as other types and arrangements of controller circuitry. Also, other types of cache entities can be used in other embodiments. For example, alternative cache entities in some embodiments can include lower level components that are part of a storage hierarchy of the storage array **505** and have local caches associated therewith.

(218) Each of the MPIO drivers **512** illustratively accesses the shared LUN **506-1** via a first path set having targets associated with the first storage controller **520-1** and via a second path set having targets associated with the second storage controller **520-2**.

(219) It is to be appreciated that different numbers, types and arrangements of hosts, storage arrays, storage controllers, caches, initiators, targets, path sets and LUNs can be used in other embodiments.

(220) As indicated above, the mapping synchronization logic **518** of the orchestrator **516** is configured to synchronize the mapping of LBA ranges to cache entities across the multiple host devices **502**, for at least the shared LUN **506-1**, and possibly additional LUNs or other logical storage devices that are shared by the multiple host devices **502**. The multiple host devices **502** are also collectively referred to herein as a host cluster. The mapping synchronization logic **518** may also be configured to handle situations in which one or more of the host devices **502** lose connectivity to one of the storage controller **520**.

(221) In some embodiments, the orchestrator **516** is more particularly implemented at least in part as an otherwise conventional PPMA suitably modified to provide the above-noted functionality. The orchestrator **516** in such an embodiment is illustratively configured to store and share the storage controller mapping information across the host devices **502**, while also performing other types of management functionality for the host devices **502**. For example, the orchestrator **516** in some embodiments is configured to collect and store the mapping information from the host devices **502**, and to share the mapping information across the host devices **502** so as to facilitate synchronization of such information between the host devices **502**. In some embodiments, this illustratively involves actively informing one or more of the host devices **502** to change the corresponding mapping in the case of a failure of a cache entity, so as to prevent workload imbalance across the storage controllers.

(222) The orchestrator **516** in some embodiments is configured to save this information in an associated database, providing for a consistent view across all the host devices **502** of the cluster.

Whenever there is a new mapping created or change in existing mapping on a given one of the host devices **502**, the corresponding MPIO driver on that host device will notify the orchestrator **516** through an event.

(223) For example, one or more of the host devices **502** may report to the orchestrator **516** that all paths to a given one of the storage controller **520** are failed. The orchestrator **516** can be configured to proactively notify all the host devices **502** of the cluster to initiate a change in their existing mapping based on at least a threshold number of host devices **502** reporting such failures.

Additionally or alternatively, the orchestrator **516** can be configured to wait for a predetermined timeout period before taking such actions, and/or may initiate tasks, possibly to be performed by one or more of the host devices **502**, to confirm that the issue is not transient.

(224) FIG. **6** shows an example of a process carried out between the orchestrator **516** and a particular MPIO driver, illustratively the MPIO driver **512-1** of the first host device **502-1**. Similar processes may be carried out between the orchestrator **516** and each of the other MPIO drivers **512** of the other host devices **502**.

(225) In this example process, it is assumed that one of the host devices **502**, namely the host device **502-1**, wants to create a new mapping entry for a given LUN. If a mapping for this LUN does not exist within MPIO software on the host device **502-1**, a new mapping entry is created, illustratively by the MPIO driver **512-1**. The orchestrator **516** collects the mapping information from the MPIO driver **512-1** of the host device **502-1**. Note that this can be an event notified from the host device **502-1** to the orchestrator **516** via the MPIO driver **512-1**. The orchestrator **516** accesses its database to check if the LUN is a shared LUN. For a shared LUN, the orchestrator **516** determines if the mapping entry created by the MPIO driver **512-1** of the host device **502-1** is inconsistent with corresponding mapping entries of other host devices **502** of the cluster, and if there is no inconsistency, the orchestrator **516** informs the MPIO driver **512-1** of the host device **502-1** to update the host entry. Otherwise, the orchestrator **516** may direct the MPIO driver **512-1** of the host device **502-1** to make a different update so as to maintain consistency in the mappings across all of the host devices **502**. The MPIO driver **512-1** of the host device **502-1** then updates the mapping entry accordingly.

(226) FIG. **7** shows another example of a process carried out between the orchestrator **516** and a particular MPIO driver, illustratively the MPIO driver **512-1** of the first host device **502-1**. Again, similar processes may be carried out between the orchestrator **516** and each of the other MPIO drivers **512** of the other host devices **502**.

(227) In this example process, if all paths to a caching entity fail, the MPIO driver **512-1** of the host device **502-1** upon detecting such a failure informs the orchestrator **516** of that condition and any observed reason for the failure. This mapping-related communication illustratively comprises a failure information packet, although other types of communications can be used. The orchestrator **516** checks if other host devices **502** have reported such a failure for the same cache entity. If at least a threshold number of host devices **502** have reported the same failure for that cache entity within a given predefined time interval, the orchestrator **516** clears its corresponding mapping entry and asks the host devices **502** to clear their respective corresponding mapping entries. This acts as an exemplary trigger for each of the host devices **502** to create a new mapping entry, and the mapping entry of the first host device **502-1** is updated as needed via the multi-path driver **512-1**. If the threshold number of host devices **502** have not reported the same failure for this cache entity within the predefined time interval, the orchestrator **516** can request for proactive testing of the paths to the cache entity to see if the failure is transient. The host device **502-1** via its MPIO driver **512-1** either clears its mapping entry or proactively tests the paths to the caching entity for the requested amount of time.

(228) In some embodiments, the orchestrator **516**, when it determines that a given cached entity is unavailable, informs all the host devices **502** in the cluster to clear their entries, so that new entries can be created for the corresponding logical blocks.

(229) These entries are illustratively saved on the orchestrator **516** and shared to the host devices **502** asynchronously, although other arrangements can be used.

(230) Illustrative embodiments avoid the sub-optimal performance that might otherwise result in situations in which clustered hosts do not interact with each other and hence can have the same logical blocks of a given shared LUN mapped to different cache entities. Moreover, illustrative embodiments provide techniques for efficiently handling situations in which hosts lose connectivity to a cache entity.

(231) The above-described cache aware features and functionality can be implemented with any storage array or other storage system which utilizes multiple cache entities that are utilized by multiple hosts to access shared LUNs or other logical storage devices.

(232) The disclosed techniques can be implemented through modification of any type of multipathing software, including Linux native multipathing software.

(233) The above-described processes, algorithms and other features and functionality disclosed herein are presented by way of illustrative example only, and other embodiments can utilize additional or alternative arrangements.

(234) Also, as mentioned previously, different instances of the above-described processes, algorithms and other synchronized mapping techniques can be performed by different MPIO drivers in different host devices, or using other types of host drivers, such as, for example, iSCSI drivers.

(235) The particular synchronized mapping arrangements described above are therefore presented by way of illustrative example only. Numerous alternative arrangements of these and other features can be used in implementing the synchronized mapping in other illustrative embodiments.

(236) As indicated previously, illustrative embodiments disclosed herein can provide a number of significant advantages relative to conventional arrangements.

(237) For example, these embodiments provide an efficient algorithm for synchronized mapping in storage cache aware load balancing.

(238) Illustrative embodiments can therefore prevent significant degradations in I/O processing performance that might otherwise occur due to excessive cache misses and slot collisions in storage cache aware load balancing.

(239) Some embodiments configure a multi-path layer of a cluster of host devices to include functionality for storage cache aware processing of I/O operations using synchronized mapping across the cluster of host devices that provides improved load balancing relative to conventional approaches.

(240) It is to be appreciated that the particular advantages described above and elsewhere herein are associated with particular illustrative embodiments and need not be present in other embodiments. Also, the particular types of information processing system features and functionality as illustrated in the drawings and described above are exemplary only, and numerous other arrangements may be used in other embodiments.

(241) It was noted above that portions of an information processing system as disclosed herein may be implemented using one or more processing platforms. Illustrative embodiments of such platforms will now be described in greater detail. These and other processing platforms may be used to implement at least portions of other information processing systems in other embodiments. A given such processing platform comprises at least one processing device comprising a processor coupled to a memory.

(242) One illustrative embodiment of a processing platform that may be used to implement at least a portion of an information processing system comprises cloud infrastructure including virtual machines implemented using a hypervisor that runs on physical infrastructure. The cloud infrastructure further comprises sets of applications running on respective ones of the virtual machines under the control of the hypervisor. It is also possible to use multiple hypervisors each providing a set of virtual machines using at least one underlying physical machine. Different sets of

virtual machines provided by one or more hypervisors may be utilized in configuring multiple instances of various components of the system.

(243) These and other types of cloud infrastructure can be used to provide what is also referred to herein as a multi-tenant environment. One or more system components such as virtual machines, or portions thereof, are illustratively implemented for use by tenants of such a multi-tenant environment.

(244) Cloud infrastructure as disclosed herein can include cloud-based systems such as Amazon Web Services, Google Cloud Platform and Microsoft Azure. Virtual machines provided in such systems can be used to implement a fast tier or other front-end tier of a multi-tier storage system in illustrative embodiments. A capacity tier or other back-end tier of such a multi-tier storage system can be implemented using one or more object stores such as Amazon S3, Google Cloud Platform Cloud Storage, and Microsoft Azure Blob Storage.

(245) In some embodiments, the cloud infrastructure additionally or alternatively comprises a plurality of containers illustratively implemented using respective operating system kernel control groups of one or more container host devices. For example, a given container of cloud infrastructure illustratively comprises a Docker container or other type of LXC implemented using a kernel control group. The containers may run on virtual machines in a multi-tenant environment, although other arrangements are possible. The containers may be utilized to implement a variety of different types of functionality within the system **100**. For example, containers can be used to implement respective compute nodes or storage nodes of a cloud-based system. Again, containers may be used in combination with other virtualization infrastructure such as virtual machines implemented using a hypervisor.

(246) Another illustrative embodiment of a processing platform that may be used to implement at least a portion of an information processing system comprises a plurality of processing devices which communicate with one another over at least one network. The network may comprise any type of network, including by way of example a global computer network such as the Internet, a WAN, a LAN, a satellite network, a telephone or cable network, a cellular network, a wireless network such as a WiFi or WiMAX network, or various portions or combinations of these and other types of networks.

(247) Each processing device of the processing platform comprises a processor coupled to a memory. The processor may comprise a microprocessor, a microcontroller, an application-specific integrated circuit (ASIC), a field-programmable gate array (FPGA), a graphics processing unit (GPU) or other type of processing circuitry, as well as portions or combinations of such circuitry elements. The memory may comprise random access memory (RAM), read-only memory (ROM), flash memory or other types of memory, in any combination. The memory and other memories disclosed herein should be viewed as illustrative examples of what are more generally referred to as “processor-readable storage media” storing executable program code of one or more software programs.

(248) Articles of manufacture comprising such processor-readable storage media are considered illustrative embodiments. A given such article of manufacture may comprise, for example, a storage array, a storage disk or an integrated circuit containing RAM, ROM, flash memory or other electronic memory, or any of a wide variety of other types of computer program products. The term “article of manufacture” as used herein should be understood to exclude transitory, propagating signals.

(249) Also included in the processing device is network interface circuitry, which is used to interface the processing device with the network and other system components, and may comprise conventional transceivers.

(250) As another example, portions of a given processing platform in some embodiments can comprise converged infrastructure.

(251) Again, these particular processing platforms are presented by way of example only, and other

embodiments may include additional or alternative processing platforms, as well as numerous distinct processing platforms in any combination, with each such platform comprising one or more computers, servers, storage devices or other processing devices.

(252) It should therefore be understood that in other embodiments different arrangements of additional or alternative elements may be used. At least a subset of these elements may be collectively implemented on a common processing platform, or each such element may be implemented on a separate processing platform.

(253) Also, numerous other arrangements of computers, servers, storage devices or other components are possible in an information processing system as disclosed herein. Such components can communicate with other elements of the information processing system over any type of network or other communication media.

(254) As indicated previously, components of an information processing system as disclosed herein can be implemented at least in part in the form of one or more software programs stored in memory and executed by a processor of a processing device. For example, at least portions of the functionality of host devices **102**, SAN **104** and storage array **105** are illustratively implemented in the form of software running on one or more processing devices. As a more particular example, the instances of path selection logic **114** and chunk mapping logic **115** may be implemented at least in part in software, as indicated previously herein.

(255) It should again be emphasized that the above-described embodiments are presented for purposes of illustration only. Many variations and other alternative embodiments may be used. For example, the disclosed techniques are applicable to a wide variety of other types of information processing systems, utilizing other arrangements of host devices, networks, storage systems, storage arrays, storage devices, processors, memories, IO queues, MPIIO drivers, path selection logic, chunk mapping logic, orchestrators, mapping synchronization logic and additional or alternative components. Also, the particular configurations of system and device elements and associated processing operations illustratively shown in the drawings can be varied in other embodiments. For example, a wide variety of different host device configurations and associated synchronized mapping arrangements can be used in other embodiments. Moreover, the various assumptions made above in the course of describing the illustrative embodiments should also be viewed as exemplary rather than as requirements or limitations. Numerous other alternative embodiments within the scope of the appended claims will be readily apparent to those skilled in the art.

Claims

1. An apparatus comprising: at least one processing device comprising a processor coupled to a memory; wherein the at least one processing device is configured: to store a mapping for each of a plurality of logical storage devices of a storage system, the mapping for a given one of the logical storage devices assigning different ranges of logical block addresses of the given logical storage device to respective different cache entities of the storage system; to receive at least one mapping-related communication from at least one of first and second host devices that have shared access to the given logical storage device; to modify the stored mapping for the given logical storage device based at least in part on the at least one received mapping-related communication; and to send at least one mapping-related communication to at least one of the first and second host devices based at least in part on the modified stored mapping; wherein receiving at least one mapping-related communication from at least one of first and second host devices that have shared access to the given logical storage device comprises receiving a notification from the first host device indicating that the first host device has generated a new mapping entry; wherein modifying of the stored mapping for the given logical storage device is controlled based at least in part on results of performance of a consistency check between the new mapping entry and one or more existing

- mapping entries of the stored mapping; and wherein the stored mapping is modified to include the new mapping entry responsive to a determination that there is no inconsistency between the new mapping entry and any corresponding existing mapping entry of the stored mapping.
2. The apparatus of claim 1 wherein the at least one processing device comprises an orchestrator device that is separate from the first and second host devices.
 3. The apparatus of claim 2 wherein the orchestrator device comprises a multi-path input-output management station coupled to each of the first and second host devices and configured to interact with first and second multi-path input-output drivers of the respective first and second host devices.
 4. The apparatus of claim 1 wherein one or more of the cache entities are each assigned multiple distinct ones of the ranges of logical block addresses.
 5. The apparatus of claim 1 wherein each of first and second host devices separately stores a version of the mapping for the given logical storage device and accesses its corresponding stored version of the mapping in conjunction with selecting paths for delivery of input-output operations to the storage system.
 6. The apparatus of claim 5 wherein selecting paths for delivery of respective input-output operations in a given one of the first and second host devices comprises: accessing the stored version of the mapping for the given logical storage device; and for each of at least a subset of the input-output operations: identifying a particular one of the cache entities based at least in part on a logical block address of the input-output operation and the stored version of the mapping; and selecting a particular one of a plurality of paths for delivery of the input-output operation to the storage system based at least in part on the identified cache entity.
 7. The apparatus of claim 1 wherein the cache entities of the storage system comprise respective components of the storage system that have respective local caches associated therewith.
 8. The apparatus of claim 7 wherein the cache entities of the storage system comprise respective storage controllers of the storage system with each such storage controller having a different local cache associated therewith.
 9. The apparatus of claim 1 wherein the at least one processing device is further configured: to determine whether or not the new mapping entry is consistent with a corresponding existing mapping entry of the stored mapping; responsive to an affirmative determination, to send the mapping-related communication to the first host device with an instruction for the first host device to utilize the new mapping entry; responsive to a negative determination, to send the mapping-related communication to the first host device with an instruction for the first host device to utilize the corresponding existing mapping entry instead of the new mapping entry.
 10. The apparatus of claim 1 wherein receiving at least one mapping-related communication from at least one of first and second host devices that have shared access to the given logical storage device comprises receiving a notification from the first host device indicating that all paths from the first host device to a particular cache entity have failed.
 11. The apparatus of claim 10 wherein the at least one processing device is further configured: to determine whether or not at least a threshold number of host devices have each reported that all paths from that host device to the particular cache entity have failed; responsive to an affirmative determination, to send the mapping-related communication to each of the first host device and one or more other host devices with an instruction for that host device to clear its mapping entry for the particular cache entity and to create a new mapping entry for the particular cache entity; responsive to a negative determination, to send the mapping-related communication to at least the first host device with an instruction for the first host device to perform one or more tests for detection of a transient failure.
 12. The apparatus of claim 1 wherein the receiving, modifying and sending are repeated to ensure that the host devices each maintain a consistent stored version of the mapping for the given logical storage device.
 13. A computer program product comprising a non-transitory processor-readable storage medium

having stored therein program code of one or more software programs, wherein the program code, when executed by at least one processing device comprising a processor coupled to a memory, causes the at least one processing device: to store a mapping for each of a plurality of logical storage devices of a storage system, the mapping for a given one of the logical storage devices assigning different ranges of logical block addresses of the given logical storage device to respective different cache entities of the storage system; to receive at least one mapping-related communication from at least one of first and second host devices that have shared access to the given logical storage device; to modify the stored mapping for the given logical storage device based at least in part on the at least one received mapping-related communication; and to send at least one mapping-related communication to at least one of the first and second host devices based at least in part on the modified stored mapping; wherein receiving at least one mapping-related communication from at least one of first and second host devices that have shared access to the given logical storage device comprises receiving a notification from the first host device indicating that the first host device has generated a new mapping entry; wherein modifying of the stored mapping for the given logical storage device is controlled based at least in part on results of performance of a consistency check between the new mapping entry and one or more existing mapping entries of the stored mapping; and wherein the stored mapping is modified to include the new mapping entry responsive to a determination that there is no inconsistency between the new mapping entry and any corresponding existing mapping entry of the stored mapping.

14. The computer program product of claim 13 wherein the at least one processing device comprises an orchestrator device that is separate from the first and second host devices.

15. The computer program product of claim 14 wherein the orchestrator device comprises a multi-path input-output management station coupled to each of the first and second host devices and configured to interact with first and second multi-path input-output drivers of the respective first and second host devices.

16. A method comprising: to store a mapping for each of a plurality of logical storage devices of a storage system, the mapping for a given one of the logical storage devices assigning different ranges of logical block addresses of the given logical storage device to respective different cache entities of the storage system; to receive at least one mapping-related communication from at least one of first and second host devices that have shared access to the given logical storage device; to modify the stored mapping for the given logical storage device based at least in part on the at least one received mapping-related communication; and to send at least one mapping-related communication to at least one of the first and second host devices based at least in part on the modified stored mapping; wherein receiving at least one mapping-related communication from at least one of first and second host devices that have shared access to the given logical storage device comprises receiving a notification from the first host device indicating that the first host device has generated a new mapping entry; wherein modifying of the stored mapping for the given logical storage device is controlled based at least in part on results of performance of a consistency check between the new mapping entry and one or more existing mapping entries of the stored mapping; and wherein the stored mapping is modified to include the new mapping entry responsive to a determination that there is no inconsistency between the new mapping entry and any corresponding existing mapping entry of the stored mapping.

17. The method of claim 16 wherein the method is performed by at least one processing device that comprises an orchestrator device that is separate from the first and second host devices.

18. The method of claim 17 wherein the orchestrator device comprises a multi-path input-output management station coupled to each of the first and second host devices and configured to interact with first and second multi-path input-output drivers of the respective first and second host devices.

19. The method of claim 16 wherein receiving at least one mapping-related communication from at least one of first and second host devices that have shared access to the given logical storage device

comprises receiving a notification from the first host device indicating that all paths from the first host device to a particular cache entity have failed.
