

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12405740
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Shah; Mansi et al.

Direct access storage for persistent services in a virtualized computing system

Abstract

An example virtualized computing system includes a cluster of hosts having a virtualization layer executing thereon and configured to manage virtual machines (VMs); a local storage device in a first host of the cluster, the local storage device configured as direct-attach storage; and an orchestration control plane, integrated with the virtualization layer, the orchestration control plane including a master server managing state of the orchestration control plane, the state including objects representing the hosts and the VMs, the orchestration control plane deploying a persistent application executing on a first VM of the VMs, the persistent application storing persistent data on the direct-attach storage.

Inventors: Shah; Mansi (San Jose, CA), Neelakantan; Subramanian (Bangalore, IN), Dickmann; Christian (Kassel, DE), Wu; Yu (Shanghai, CN)

Applicant: VMware, Inc. (Palo Alto, CA)

Family ID: 80625773

Assignee: VMware LLC (Palo Alto, CA)

Appl. No.: 18/339923

Filed: June 22, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20230333765 A1	Oct. 19, 2023

Foreign Application Priority Data

WO	PCT/CN2020/114691	Sep. 11, 2020
----	-------------------	---------------

Related U.S. Application Data

continuation parent-doc US 17081793 20201027 US 11720271 child-doc US 18339923

Publication Classification

Int. Cl.: G06F3/06 (20060101); G06F9/455 (20180101); G06F9/54 (20060101)

U.S. Cl.:

CPC G06F3/0644 (20130101); G06F3/0604 (20130101); G06F3/0664 (20130101); G06F3/067 (20130101); G06F9/45558 (20130101); G06F9/547 (20130101); G06F2009/45579 (20130101)

Field of Classification Search

CPC: G06F (3/0644); G06F (3/0604); G06F (3/0664); G06F (3/067); G06F (9/547); G06F (2009/45579); G06F (9/45558)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2015/0378853	12/2014	Sevigny et al.	N/A	N/A
2016/0048408	12/2015	Madhu et al.	N/A	N/A
2016/0162338	12/2015	Sathyamurthy	718/1	G06F 9/45558
2017/0054603	12/2016	Kulkarni et al.	N/A	N/A
2018/0101395	12/2017	Aleksandrov	N/A	G06F 9/5077
2019/0227845	12/2018	Sridhar	N/A	G06F 9/5044
2020/0034240	12/2019	Natanzon	N/A	G06F 11/1448
2020/0042392	12/2019	Alluboyina et al.	N/A	N/A
2020/0059401	12/2019	Holmes	N/A	H04L 41/0803
2020/0250133	12/2019	Bono et al.	N/A	N/A
2020/0371698	12/2019	Chen	N/A	G06F 3/0629
2021/0149774	12/2020	Denneman et al.	N/A	N/A
2021/0168046	12/2020	Madala	N/A	H04L 41/20
2021/0279187	12/2020	Puder	N/A	G06F 13/1668
2021/0294774	12/2020	Keller	N/A	G06F 16/125
2021/0294775	12/2020	Keller	N/A	G06F 16/1734
2021/0303160	12/2020	Lieblich	N/A	G06F 3/067
2021/0303169	12/2020	Tagar	N/A	G06F 3/0608
2021/0303401	12/2020	Yarimi	N/A	G06F 11/2094
2021/0303407	12/2020	Keller	N/A	G06F 16/9027
2021/0303480	12/2020	Keller	N/A	G06F 12/0824
2021/0373796	12/2020	Matosevich	N/A	G06F 3/0689
2021/0409349	12/2020	Grunwald et al.	N/A	N/A
2022/0004320	12/2021	Matosevich	N/A	G06F 3/0659
2022/0050932	12/2021	Levy	N/A	G06F 3/0673

Background/Summary

CROSS-REFERENCE (1) This application is a continuation of U.S. patent application Ser. No. 17/081,793, filed Oct. 27, 2020, which is based upon and claims the benefit of priority from International Patent Application No. PCT/CN2020/114691, filed on Sep. 11, 2020, the entire contents of which are incorporated herein by reference.

BACKGROUND

(1) Applications today are deployed onto a combination of virtual machines (VMs), containers, application services, and more implemented in a virtualized computing system. For deploying such applications, a container orchestrator (CO) known as Kubernetes® has gained in popularity among application developers. Kubernetes provides a platform for automating deployment, scaling, and operations of application containers across clusters of hosts. It offers flexibility in application development and offers several useful tools for scaling.

(2) Some applications are stateful in that they store their data persistently in storage provided by the virtualized computing system (referred to herein as “stateful applications” or “persistent services”). For example, the virtualized computing system can provide a storage area network (SAN) that implements shared storage for various stateful applications to store their persistent data. Another type of shared storage is a virtual SAN (vSAN), which is an aggregation of local storage devices in hosts into shared storage for use by all hosts. A vSAN can be a policy-based datastore, meaning each object created therein can specify a level of replication and protection. The vSAN can then perform automatic placement decisions for replicas across the local storage devices of the hosts. A SAN can function similarly to a vSAN but includes storage devices separate from the hosts and accessible through a network. A stateful application deployed by Kubernetes, for example, can have one or more persistent volumes (PVs) stored in shared storage (e.g., a SAN or a vSAN), where the PVs enjoy the replication and protection provided by the underlying shared storage.

(3) Stateful applications (e.g., databases) can provide their own replication and data protection services for their data. Such stateful applications do not benefit from paying the cost of obtaining the same data protection services from the underlying shared storage system. For example, if a stateful application performs 3-way replication of its data, and the underlying storage system performs 2-way replication of the data, then six copies of the same data are created. However, the stateful application only obtains fault tolerance of the underlying storage system (e.g., the 2-way replication). The drop in fault tolerance occurs because the underlying shared storage system is not aware of the correlation between application replicas and is free to use the exact same placement for all replicas created by the application.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1 is a block diagram of a virtualized computing system in which embodiments described herein may be implemented.

(2) FIG. 2 is a block diagram depicting a software platform according an embodiment.

(3) FIG. 3 is a block diagram of a supervisor Kubernetes master according to an embodiment.

(4) FIG. 4 is a block diagram depicting a persistent application deployed on a supervisor cluster

according to an embodiment,

(5) FIG. 5 is a block diagram depicting deployment of a persistent application across a plurality of hosts and using direct-attach storage according to an embodiment.

(6) FIG. 6 is a flow diagram depicting a method of deploying a persistent application with direct-attached storage according to an embodiment.

DETAILED DESCRIPTION

(7) Techniques for providing a paravirtual storage layer for a container orchestrator in a virtualized computing system are described. In embodiments described herein, the virtualized computing system includes a cluster of physical servers (“hosts”) referred to as a “host cluster.” The host cluster includes a virtualization layer, executing on hardware platforms of the hosts, which supports execution of virtual machines (VMs). A virtualization management server manages the host cluster, the virtualization layer, and the VMs executing thereon. The virtualized computing system includes shared storage accessible by the host cluster. The container orchestrator executes in the virtualized computing system (e.g., on one or more VMs) and is configured to deploy and manage applications in the host cluster. In embodiments, the container orchestrator is a Kubernetes system that deploys and manages containerized applications in a cluster of VMs (a “Kubernetes cluster”).

(8) In one or more embodiments, the orchestration control plane comprises a supervisor container orchestrator having extensions that cooperate with the virtualization management server and agents installed in the virtualization layer. A host cluster having the orchestration control plane is referred to herein as a “supervisor cluster.” A user interacts with the orchestration control plane to deploy and manage applications executing on the supervisor cluster. In embodiments, the orchestration control plane uses hosts to implement nodes, and VMs to implement pods, of a Kubernetes cluster. Kubernetes pods are implemented as “pod VMs,” each of which includes a kernel and container engine that supports execution of containers. The container orchestrator (e.g., Kubernetes) executes in VMs alongside the pod VMs.

(9) In one or more embodiments, local storage devices in a host are excluded from use as a virtual storage area network (vSAN) and are instead configured as direct-attach storage. A persistent application deployed by the orchestration control plane can use direct-attach storage for placement of persistent data, rather than on the vSAN. This obviates the need to traverse the vSAN software stack, allowing the persistent application direct access to the local storage devices comprising the direct-attach storage. The persistent application controls placement of the persistent data across the direct-attach storage devices. These and further advantages and aspects of the disclosed techniques are described below with respect to the drawings.

(10) FIG. 1 is a block diagram of a virtualized computing system **100** in which embodiments described herein may be implemented. System **100** includes a cluster of hosts **120** (“host cluster **118**”) that may be constructed on server-grade hardware platforms such as an x86 architecture platforms. For purposes of clarity, only one host cluster **118** is shown. However, virtualized computing system **100** can include many of such host clusters **118**. As shown, a hardware platform **122** of each host **120** includes conventional components of a computing device, such as one or more central processing units (CPUs) **160**, system memory (e.g., random access memory (RAM) **162**), one or more network interface controllers (NICs) **164**, and optionally local storage **163**. CPUs **160** are configured to execute instructions, for example, executable instructions that perform one or more operations described herein, which may be stored in RAM **162**. NICs **164** enable host **120** to communicate with other devices through a physical network **180**. Physical network **180** enables communication between hosts **120** and between other components and hosts **120** (other components discussed further herein). Physical network **180** can include a plurality of VLANs to provide external network virtualization as described further herein.

(11) In the embodiment illustrated in FIG. 1, hosts **120** access shared storage **170** by using NICs **164** to connect to network **180**. In another embodiment, each host **120** contains a host bus adapter (HBA) through which input/output operations (**10s**) are sent to shared storage **170** over a separate

network (e.g., a fibre channel (FC) network). Shared storage **170** include one or more storage arrays, such as a storage area network (SAN), network attached storage (NAS), or the like. Shared storage **170** may comprise magnetic disks, solid-state disks, flash memory, and the like as well as combinations thereof. In some embodiments, hosts **120** include local storage **163** (e.g., hard disk drives, solid-state drives, etc.). Local storage **163** in each host **120** can be aggregated and provisioned as part of a virtual SAN (vSAN), which is another form of shared storage **170**. Virtualization management server **116** can select which local storage devices in hosts **120** are part of a vSAN for host cluster **118**. In further embodiments described herein, virtualization management server **116** can omit some local storage devices from the vSAN for host cluster **118** for use as direct-attach storage. A user can request direct-attach storage for an application, as described further herein.

(12) A software platform **124** of each host **120** provides a virtualization layer, referred to herein as a hypervisor **150**, which directly executes on hardware platform **122**. In an embodiment, there is no intervening software, such as a host operating system (OS), between hypervisor **150** and hardware platform **122**. Thus, hypervisor **150** is a Type-1 hypervisor (also known as a “bare-metal” hypervisor). As a result, the virtualization layer in host cluster **118** (collectively hypervisors **150**) is a bare-metal virtualization layer executing directly on host hardware platforms. Hypervisor **150** abstracts processor, memory, storage, and network resources of hardware platform **122** to provide a virtual machine execution space within which multiple virtual machines (VM) may be concurrently instantiated and executed. One example of hypervisor **150** that may be configured and used in embodiments described herein is a VMware ESXi™ hypervisor provided as part of the VMware vSphere® solution made commercially available by VMware, Inc. of Palo Alto, CA.

(13) In the example of FIG. 1, host cluster **118** is enabled as a “supervisor cluster,” described further herein, and thus VMs executing on each host **120** include pod VMs **130** and native VMs **140**. A pod VM **130** is a virtual machine that includes a kernel and container engine that supports execution of containers, as well as an agent (referred to as a pod VM agent) that cooperates with a controller of an orchestration control plane **115** executing in hypervisor **150** (referred to as a pod VM controller). An example of pod VM **130** is described further below with respect to FIG. 2. VMs **130/140** support applications **141** deployed onto host cluster **118**, which can include containerized applications (e.g., executing in either pod VMs **130** or native VMs **140**) and applications executing directly, on guest operating systems (non-containerized) (e.g., executing in native VMs **140**). One specific application discussed further herein is a guest cluster executing as a virtual extension of a supervisor cluster. Some VMs **130/140**, shown as support VMs **145**, have specific functions within host cluster **118**. For example, support VMs **145** can provide control plane functions, edge transport functions, and the like. An embodiment of software platform **124** is discussed further below with respect to FIG. 2.

(14) Host cluster **118** is configured with a software-defined (SD) network layer **175**. SD network layer **175** includes logical network services executing on virtualized infrastructure in host cluster **118**. The virtualized infrastructure that supports the logical network services includes hypervisor-based components, such as resource pools, distributed switches, distributed switch port groups and uplinks, etc., as well as VM-based components, such as router control VMs, load balancer VMs, edge service VMs, etc. Logical network services include logical switches, logical routers, logical firewalls, logical virtual private networks (VPNs), logical load balancers, and the like, implemented on top of the virtualized infrastructure. In embodiments, virtualized computing system **100** includes edge transport nodes **178** that provide an interface of host cluster **118** to an external network (e.g., a corporate network, the public Internet, etc.). Edge transport nodes **178** can include a gateway between the internal logical networking of host cluster **118** and the external network. Edge transport nodes **178** can be physical servers or VMs. For example, edge transport nodes **178** can be implemented in support VMs **145** and include a gateway of SD network layer **175**. Various clients **119** can access service(s) in virtualized computing system through edge transport nodes **178**.

(including VM management client **106** and Kubernetes client **102**, which as logically shown as being separate by way of example).

(15) Virtualization management server **116** is a physical or virtual server that manages host cluster **118** and the virtualization layer therein. Virtualization management server **116** installs agent(s) **152** in hypervisor **150** to add a host **120** as a managed entity. Virtualization management server **116** logically groups hosts **120** into host cluster **118** to provide cluster-level functions to hosts **120**, such as VM migration between hosts **120** (e.g., for load balancing), distributed power management, dynamic VM placement according to affinity and anti-affinity rules, and high-availability. The number of hosts **120** in host cluster **118** may be one or many. Virtualization management server **116** can manage more than one host cluster **118**.

(16) In an embodiment, virtualization management server **116** further enables host cluster **118** as a supervisor cluster **101**. Virtualization management server **116** installs additional agents **152** in hypervisor **150** to add host **120** to supervisor cluster **101**. Supervisor cluster **101** integrates an orchestration control plane **115** with host cluster **118**. In embodiments, orchestration control plane **115** includes software components that support a container orchestrator, such as Kubernetes, to deploy and manage applications on host cluster **118**. By way of example, a Kubernetes container orchestrator is described herein. In supervisor cluster **101**, hosts **120** become nodes of a Kubernetes cluster and pod VMs **130** executing on hosts **120** implement Kubernetes pods. Orchestration control plane **115** includes supervisor Kubernetes master **104** and agents **152** executing in virtualization layer (e.g., hypervisors **150**). Supervisor Kubernetes master **104** includes control plane components of Kubernetes, as well as custom controllers, custom plugins, scheduler extender, and the like that extend Kubernetes to interface with virtualization management server **116** and the virtualization layer. For purposes of clarity, supervisor Kubernetes master **104** is shown as a separate logical entity. For practical implementations, supervisor Kubernetes master **104** is implemented as one or more VM(s) **130/140** in host cluster **118**. Further, although only one supervisor Kubernetes master **104** is shown, supervisor cluster **101** can include more than one supervisor Kubernetes master **104** in a logical cluster for redundancy and load balancing.

(17) In an embodiment, virtualized computing system **100** further includes a storage service **110** that implements a storage provider in virtualized computing system **100** for container orchestrators. In embodiments, storage service **110** manages lifecycles of storage volumes (e.g., virtual disks) that back persistent volumes used by containerized applications executing in host cluster **118**. A container orchestrator such as Kubernetes cooperates with storage service **110** to provide persistent storage for the deployed applications. In the embodiment of FIG. 1, supervisor Kubernetes master **104** cooperates with storage service **110** to deploy and manage persistent storage in the supervisor cluster environment. Other embodiments described below include a vanilla container orchestrator environment and a guest cluster environment. Storage service **110** can execute in virtualization management server **116** as shown or operate independently from virtualization management server **116** (e.g., as an independent physical or virtual server).

(18) In an embodiment, virtualized computing system **100** further includes a network manager **112**. Network manager **112** is a physical or virtual server that orchestrates SD network layer **175**. In an embodiment, network manager **112** comprises one or more virtual servers deployed as VMs. Network manager **112** installs additional agents **152** in hypervisor **150** to add a host **120** as a managed entity, referred to as a transport node. In this manner, host cluster **118** can be a cluster **103** of transport nodes. One example of an SD networking platform that can be configured and used in embodiments described herein as network manager **112** and SD network layer **175** is a VMware MX® platform made commercially available by VMware, Inc. of Palo Alto, CA.

(19) Network manager **112** can deploy one or more transport zones in virtualized computing system **100**, including VLAN transport zone(s) and an overlay transport zone. A VLAN transport zone spans a set of hosts **120** (e.g., host cluster **118**) and is backed by external network virtualization of physical network **180** (e.g., a VLAN). One example VLAN transport zone uses a management

VLAN **182** on physical network **180** that enables a management network connecting hosts **120** and the VI control plane (e.g., virtualization management server **116** and network manager **112**). An overlay transport zone using overlay VLAN **184** on physical network **180** enables an overlay network that spans a set of hosts **120** (e.g., host cluster **118**) and provides internal network virtualization using software components (e.g., the virtualization layer and services executing in VMs). Host-to-host traffic for the overlay transport zone is carried by physical network **180** on the overlay VLAN **184** using layer-2-over-layer-3 tunnels. Network manager **112** can configure SD network layer **175** to provide a cluster network **186** using the overlay network. The overlay transport zone can be extended into at least one of edge transport nodes **178** to provide ingress/egress between cluster network **186** and an external network.

(20) In an embodiment, system **100** further includes an image registry **190**. As described herein, containers of supervisor cluster **101** execute in pod VMs **130**. The containers in pod VMs **130** are spun up from container images managed by image registry **190**. Image registry **190** manages images and image repositories for use in supplying images for containerized applications.

(21) Virtualization management server **116** and network manager **112** comprise a virtual infrastructure (VI) control plane **113** of virtualized computing system **100**. Virtualization management server **116** can include a supervisor cluster service **109**, storage service **110**, and VI services **108**. Supervisor cluster service **109** enables host cluster **118** as supervisor cluster **101** and deploys the components of orchestration control plane **115**. VI services **108** include various virtualization management services, such as a distributed resource scheduler (DRS), high-availability (HA) service, single sign-on (SSO) service, virtualization management daemon, and the like. DRS is configured to aggregate the resources of host cluster **118** to provide resource pools and enforce resource allocation policies. DRS also provides resource management in the form of load balancing, power management, VM placement, and the like. HA service is configured to pool VMs and hosts into a monitored cluster and, in the event of a failure, restart VMs on alternate hosts in the cluster. A single host is elected as a master, which communicates with the HA service and monitors the state of protected VMs on subordinate hosts. The HA service uses admission control to ensure enough resources are reserved in the cluster for VM recovery when a host fails. SSO service comprises security token service, administration server, directory service, identity management service, and the like configured to implement an SSO platform for authenticating users. The virtualization management daemon is configured to manage objects, such as data centers, clusters, hosts, VMs, resource pools, datastores, and the like.

(22) A VI admin can interact with virtualization management server **116** through a VM management client **106**. Through VM management client **106**, a VI admin commands virtualization management server **116** to form host cluster **118**, configure resource pools, resource allocation policies, and other cluster-level functions, configure storage and networking, enable supervisor cluster **101**, deploy and manage image registry **190**, and the like.

(23) Kubernetes client **102** represents an input interface for a user to supervisor Kubernetes master **104**. For example, Kubernetes client **102** can be kubectl. Through Kubernetes client **102**, a user submits desired states of the Kubernetes system, e.g., as YAML documents, to supervisor Kubernetes master **104**. In embodiments, the user submits the desired states within the scope of a supervisor namespace. A “supervisor namespace” is a shared abstraction between VI control plane **113** and orchestration control plane **115**. Each supervisor namespace provides resource-constrained and authorization-constrained units of multi-tenancy. A supervisor namespace provides resource constraints, user-access constraints, and policies (e.g., storage policies, network policies, etc.). Resource constraints can be expressed as quotas, limits, and the like with respect to compute (CPU and memory), storage, and networking of the virtualized infrastructure (host cluster **118**, shared storage **170**, SD network layer **175**). User-access constraints include definitions of users, roles, permissions, bindings of roles to users, and the like. Each supervisor namespace is expressed within orchestration control plane **115** using a namespace native to orchestration control plane **115**.

(e.g., a Kubernetes namespace or generally a “native namespace”), which allows users to deploy applications in supervisor cluster **101** within the scope of supervisor namespaces. In this manner, the user interacts with supervisor Kubernetes master **104** to deploy applications in supervisor cluster **101** within defined supervisor namespaces.

(24) While FIG. 1 shows an example of a supervisor cluster **101**, the techniques described herein do not require a supervisor cluster **101**. In some embodiments, host cluster **118** is not enabled as a supervisor cluster **101**. In such case, supervisor Kubernetes master **104**, Kubernetes client **102**, pod VMs **130**, supervisor cluster service **109**, and image registry **190** can be omitted. While host cluster **118** is shown as being enabled as a transport node cluster **103**, in other embodiments network manager **112** can be omitted. In such case, virtualization management server **116** functions to configure SD network layer **175**.

(25) FIG. 2 is a block diagram depicting software platform **124** according to an embodiment. As described above, software platform **124** of host **120** includes hypervisor **150** that supports execution of VMs, such as pod VMs **130**, native VMs **140**, and support VMs **145**. In an embodiment, hypervisor **150** includes a VM management daemon **213**, a host daemon **214**, a pod VM controller **216**, an image service **218**, and network agents **222**. VM management daemon **213** is an agent **152** installed by virtualization management server **116**. VM management daemon **213** provides an interface to host daemon **214** for virtualization management server **116**. Host daemon **214** is configured to create, configure, and remove VMs (e.g., pod VMs **130** and native VMs **140**).

(26) Pod VM controller **216** is an agent **152** of orchestration control plane **115** for supervisor cluster **101** and allows supervisor Kubernetes master **104** to interact with hypervisor **150**. Pod VM controller **216** configures the respective host as a node in supervisor cluster **101**. Pod VM controller **216** manages the lifecycle of pod VMs **130**, such as determining when to spin-up or delete a pod VM. Pod VM controller **216** also ensures that any pod dependencies, such as container images, networks, and volumes are available and correctly configured. Pod VM controller **216** is omitted if host cluster **118** is not enabled as a supervisor cluster **101**.

(27) Image service **218** is configured to pull container images from image registry **190** and store them in shared storage **170** such that the container images can be mounted by pod VMs **130**. Image service **218** is also responsible for managing the storage available for container images within shared storage **170**. This includes managing authentication with image registry **190**, assuring providence of container images by verifying signatures, updating container images when necessary, and garbage collecting unused container images. Image service **218** communicates with pod VM controller **216** during spin-up and configuration of pod VMs **130**. In some embodiments, image service **218** is part of pod VM controller **216**. In other embodiments, image service **218** utilizes system VMs **130/140** in support VMs **145** to fetch images, convert images to container image virtual disks, and cache container image virtual disks in shared storage **170**.

(28) Network agents **222** comprises agents **152** installed by network manager **112**. Network agents **222** are configured to cooperate with network manager **112** to implement logical network services. Network agents **222** configure the respective host as a transport node in a cluster **103** of transport nodes.

(29) Each pod VM **130** has one or more containers **206** running therein in an execution space managed by container engine **208**. The lifecycle of containers **206** is managed by pod VM agent **212**. Both container engine **208** and pod VM agent **212** execute on top of a kernel **210** (e.g., a Linux® kernel). Each native VM **140** has applications **202** running therein on top of an OS **204**. Native VMs **140** do not include pod VM agents and are isolated from pod VM controller **216**. Container engine **208** can be an industry-standard container engine, such as libcontainer, runc, or containerd. Pod VMs **130**, pod VM controller **216**, and image service **218** are omitted if host cluster **118** is not enabled as a supervisor cluster **101**.

(30) FIG. 3 is a block diagram of supervisor Kubernetes master **104** according to an embodiment. Supervisor Kubernetes master **104** includes application programming interface (API) server **302**, a

state database **303**, a scheduler **304**, controllers **308**, and plugins **319**. Controllers **308** can include, for example, standard Kubernetes controllers, as well as custom controllers, such as a VM controller, guest cluster controllers, and platform lifecycle controller (PLC). Plugins **319** can include, for example, a network plugin and a storage plugin. In the embodiment shown, controllers **308** include persistent services controller **320** and one or more custom service controllers **322**. Plugins **319** include storage plugin **324**.

(31) API server **302** provides an API for use by Kubernetes client **102** (e.g., kube-api server). API server **302** is the front end of orchestration control plane **115**. The Kubernetes API provides a declarative schema for creating, updating, deleting, and viewing objects. State database **303** stores the state of supervisor cluster **101** (e.g., etcd) as objects created by API server **302**. A user can provide application specification data to API server **302** that defines various objects supported by the API (e.g., as a YAML document). The objects have specifications that represent the desired state. State database **303** stores the objects defined by application specification data as part of the supervisor cluster state.

(32) Namespaces provide scope for Kubernetes objects. Namespaces are objects themselves maintained in state database **303**. A namespace can include resource quotas, limit ranges, role bindings, and the like that are applied to objects declared within its scope. A VI admin can cooperate with VM management server **116** to define supervisor namespaces for supervisor cluster **101**. A supervisor namespace is a resource-constrained and authorization-constrained unit of multi-tenancy managed by VM management server **116**. State database **303** stores namespace objects associated with the supervisor namespaces. VM management server **116** creates a namespace object in supervisor Kubernetes master **104** for each supervisor namespace, pushing down resource constraints and authorization constraints into orchestration control plane **115**. A namespace is an example of a standard Kubernetes object. State database **303** can store various Kubernetes objects **340**, including namespaces.

(33) Scheduler **304** watches state database **303** for newly created pods with no assigned node. A pod is an object supported by API server **302** that is a group of one or more containers, with network and storage, and a specification on how to execute. Scheduler **304** selects candidate nodes in supervisor cluster **101** for pods. Scheduler **304** cooperates with scheduler extender **306**, which interfaces with VM management server **116**. Scheduler extender **306** cooperates with VM management server **116** such as with resource scheduler **108**) to select nodes from candidate sets of nodes and provide identities of hosts **120** corresponding to the selected nodes. For each pod, scheduler **304** also converts the pod specification to a pod VM specification, and scheduler extender **306** asks VM management server **116** to reserve a pod VM on the selected host **120**. Scheduler **304** updates pods in state database **303** with host identifiers.

(34) A controller **308** tracks objects in state database **303** of at least one resource type. Controller(s) **308** are responsible for making the current state of supervisor cluster **101** come closer to the desired state as stored in state database **303**. A controller **308** can carry out action(s) by itself, send messages to API server **302** to have side effects, and/or interact with external systems. A PLC, for example, is responsible for tracking pods that have assigned nodes without pod VM identifiers. The PLC cooperates with VM management server **116** to commit reserved pod VMs for pods. VM management server **116** returns a pod VM identifier to the PLC, which in turn updates the pod in state database **303**.

(35) Pods are native objects of Kubernetes. The Kubernetes API can be extended with custom APIs **305** to allow orchestration and management of custom objects **307**. A custom resource definition (CRD) can be used to define a custom object **307** to be handled by API server **302**. Alternatively, an extension API server can be used to introduce a custom object **307** by API server aggregation, where the extension API server is fully responsible for the custom resource. A user interacts with custom APIs **305** of API server **302** to create custom objects **307** tracked in state database **303**. A controller **308** is used to watch for and actuate on custom objects **307** declared in state database

303.

(36) In an embodiment, orchestration control plane **115** is extended to support orchestration of persistent services. This extensibility can be implemented using either CRDs or an extension API server in supervisor Kubernetes master **104**. A user or a controller **308** can invoke a custom API to create supervisor service objects, which represent custom persistent services. Persistent services controller **320** monitors for supervisor service objects and creates instances of custom persistent service controllers **322**. Each custom persistent service controller **322** manages the lifecycle of a custom persistent service (e.g., a database).

(37) For example, a user can specify a supervisor service object for a database application. In turn, persistent services controller **320** creates an instance of custom persistent service controller **322** for the database application. Thereafter, users can create instances of the database application through specification of database application CRDs monitored by custom persistent service controller **322**.

(38) Plugins **319** provide a well-defined interface to replace a set of functionality of the Kubernetes control plane. A network plugin is responsible for configuration of logical networking of SD networking **175** to satisfy the needs of network-related resources. Storage plugin **324** is responsible for providing a standardized interface for persistent storage lifecycle and management to satisfy the needs of resources requiring persistent storage. Storage plugin **324** provides an interface between orchestration control plane **115** and storage service **110**.

(39) FIG. 4 is a block diagram depicting a persistent application deployed on a supervisor cluster according to an embodiment. Supervisor cluster **1111** implemented by a software-defined data center (SDDC) **402**. SDDC **402** includes virtualized computing system **100** shown in FIG. 1, including host cluster **118**, VM management server **116**, network manager **112**, storage manager **110**, shared storage **170**, and SD networking **175**. SDDC **402** includes VI control plane **113** for managing a virtualization layer of host cluster **118**, along with shared storage **170** and SD networking **175**. In particular, SDDC **402** includes a vSAN **406** and direct-attach storage **408**. The vSAN **406** is implemented using local storage **163** from a plurality of hosts **120**. Direct-attach storage **408** includes a plurality of storage devices in local storage **163** of a plurality of hosts **120**. Direct-attach storage **408** is exclusive of vSAN **406** (e.g., the storage devices of direct-attach storage **408** are not part of vSAN **406**). A VI admin interacts with VM management server **116** (and optionally network manager **112**) of VI control plane **113** to configure SDDC **402** to implement supervisor cluster **101**.

(40) Supervisor cluster **101** includes orchestration control plane **115**, which includes supervisor Kubernetes master(s) **104** and pod VM controllers **216**. The VI admin interacts with VM management server **116** to create supervisor namespaces **417**. Each supervisor namespace **417** includes a resource pool and authorization constraints. The resource pool includes various resource constraints on supervisor namespace **417** (e.g., reservation, limits, and share (RLS) constraints). Authorization constraints provide for which roles are permitted to perform which operations in supervisor namespace **417** (e.g., allowing VI admin to create, manage access, allocate resources, view, and create objects; allowing users to view and create objects; etc.). A user interacts with Kubernetes master **104** to deploy applications on supervisor cluster **101** within scopes of supervisor namespaces **417**. In the example, a user deploys an application **423** on pod VM(s) **130**, an application **426** on native VM(s) **140**, a persistent application **428** on pod VM(s) **130** and/or native VM(s) **140**.

(41) Each supervisor namespace **417** also includes storage classes **404** to be used when creating persistent volumes. In embodiments, one or more storage classes **404** are associated with vSAN **406**. Persistent volume claims for applications (e.g., applications **423**, **426**) can reference storage classes **404** for vSAN **406**. In such case, orchestration control plane **115** (through storage plugin **324**) satisfies those persistent volume claims with persistent volumes stored in vSAN **406**. In embodiments, one or more storage classes **404** are associated with direct-attach storage **408**. Persistent volume claims for applications (e.g., persistent application **428**) can reference storage

classes **404** for direct-attach storage **408**. In such case, orchestration control plane **115** (through storage plugin **324**) satisfies those persistent volume claims with persistent volumes stored in direct-attach storage **408**. In embodiments, storage classes **404** can be differentiated based on performance, where each storage class **404** can be implemented using a different type of storage (e.g., NVMe, SSD, hard disks, etc.). In embodiments, each storage class defined for direct-attach storage is an abstraction of the underlying local storage devices and their properties. Each storage class for the direct-attach storage is agnostic of the underlying local storage devices.

(42) In embodiments, direct-attach storage **408** includes local storage devices in a plurality of hosts **120**. These local storage devices are exempted from being used with vSAN **406**. Each individual local storage device (e.g., NVMe devices, SSD devices, hard disk devices, etc.) presents a separate datastore available for placement of persistent volumes. Storage devices of similar performance characteristics can be grouped together into separate tiers and assigned to separate storage classes **404**. In embodiments, VI admin can interact with virtualization management server **116** to manage each tier of direct-attach storage **408**, which obviates the need for the VI admin to manage each storage device individually. Thus, persistent applications can obtain direct access to local storage devices in direct-attach storage **408**, while VI admin can manage collections of local storage devices by tier for ease of management.

(43) FIG. 5 is a block diagram depicting deployment of a persistent application across a plurality of hosts and using direct-attach storage according to an embodiment. As shown in FIG. 5, persistent applications **428-1** and **428-2** are instances of the persistent application. Orchestration control plane **115** places persistent application **428-1** in a host **120-1** and persistent application **428-2** in a host **120-2**. As such, the persistent application can be replicated across multiple hosts **120** for redundancy. In the embodiment, the user specified the persistent application to have two persistent volume claims (PVCs) with a direct-attach storage class (assume one tier for clarity). In an embodiment, in addition to specifying the direct-attach storage class, the user can also specify anti-affinity rules for PVs so that they are placed on separate local storage devices. In the embodiment, orchestration control plane **115** binds two PVs **506A** and **506B** to satisfy the two PVCs, respectively. Orchestration control plane **115** places PV **506A** on a storage device **502-1** of host **120-1**, and PV **506B** on a storage device **502-2** of host **120-1**. Persistent application **428-1** accesses storage devices **502-1** and **502-2** directly without passing through a vSAN or other storage orchestrator. Likewise, orchestration control plane **115** binds two PVs **508A** and **508B** to the two PVCs for persistent application **428-2**. Orchestration control plane **115** places PV **508A** on a storage device **504-1** and PV **508B** on a storage device **504-2**, each in host **120-2**. In the embodiment, PV **506A** and PV **508A** store two replicas of first persistent data, and PV **506B** and **508B** store two replicas of second persistent data.

(44) FIG. 6 is a flow diagram depicting a method **600** of deploying a persistent application with direct-attached storage according to an embodiment. Method **600** can be performed by software executing in supervisor cluster **101**, which comprises software executing on CPU, memory, storage, and network resources managed by a virtualization layer (e.g., a hypervisor).

(45) Method **600** begins at step **602**, where an administrator (e.g., VI Admin) adds a persistent service to supervisor cluster **101** within a supervisor namespace. In an embodiment, admin interacts with supervisor Kubernetes master **104** to specify the persistent service using the custom CRD(s) monitored by persistent services controller **320**. In response, persistent services controller **320** generates an instance of custom persistent service controller **322** for the persistent service. At step **604**, the administrator selects local storage devices in hosts **120** for direct attach storage and creates storage classes in the supervisor namespace. For example, VI admin can interact with virtualization management server **116** to select different types of storage devices in hosts **120** for use as direct-attach storage, such as NVMe, SSD, and hard disks. Virtualization management server **116** marks the selected local storage devices for direct-attach storage and exempts them from any other orchestrated storage for the cluster (e.g., vSAN). Virtualization management server **116** causes each

direct-attach storage device to be formatted and to present a pool of storage (a datastore) for placement of PVs by orchestration control plane **115**.

(46) At step **606**, a user creates a persistent application as an instance of the persistent service. In an embodiment, the user interacts with supervisor Kubernetes master **104** to specify the persistent application for implementation in pods. For example, at step **608**, the user can specify a replica count for the persistent service. Each replica of the persistent service will execute in a different pod on a different node (host **120**). At step **610**, the user can specify one or more PVCs for the persistent application. Each PVC includes a specified storage class for direct-attach storage. Optionally, at step **612**, the user can specify anti-affinity rules for the PVCs. For example, as shown in the example of FIG. 5, the persistent application can include two PVCs. Instances of PVs bound to the two PVCs are stored on separate storage devices within each host. In general, custom persistent service controller **322** can use annotations defined by API server **302** to specify how the persistent data is stored on the direct-attach storage. For example, the annotations can determine placement of the persistent data across a plurality of local storage devices configured as direct attach storage. Thus, a persistent application can obtain better performance by spreading the persistent data wider across the local storage devices or better fault tolerance by ensuring that two redundant shards of persistent data are not co-located on the same local storage device.

(47) At step **614**, orchestration control plane **115** provisions pods across nodes (hosts **120**) for the persistent application based on the replica count. For example, if the replica count is three, orchestration control plane **115** provisions three pods across three separate hosts **120**. At step **616**, orchestration control plane **115** binds PVs to the PVC(s) and places the PVs on direct-attach storage local to each node (host **120**) based on assigned storage class and anti-affinity rules (if any). For each PVC, orchestration control plane (e.g., supervisor Kubernetes master **104**) binds a PV thereto and places the PV on a direct-attach storage device based on the selected storage class. Within each host **120**, the PV is mounted within the pod executing the replica of the persistent application. In cases where the persistent application includes multiple PVCs with anti-affinity rules, orchestration control plane **115** places PVs on different direct-attach storage devices within each host **120** to satisfy the anti-affinity rules.

(48) In method **600**, steps **602** and **604** can be performed by the administrator once initially. Thereafter, steps **606-616** can be performed each time a user creates an instance of the persistent service.

(49) One or more embodiments of the invention also relate to a device or an apparatus for performing these operations. The apparatus may be specially constructed for required purposes, or the apparatus may be a general-purpose computer selectively activated or configured by a computer program stored in the computer. Various general-purpose machines may be used with computer programs written in accordance with the teachings herein, or it may be more convenient to construct a more specialized apparatus to perform the required operations.

(50) The embodiments described herein may be practiced with other computer system configurations including hand-held devices, microprocessor systems, microprocessor-based or programmable consumer electronics, minicomputers, mainframe computers, etc.

(51) One or more embodiments of the present invention may be implemented as one or more computer programs or as one or more computer program modules embodied in computer readable media. The term computer readable medium refers to any data storage device that can store data which can thereafter be input to a computer system. Computer readable media may be based on any existing or subsequently developed technology that embodies computer programs in a manner that enables a computer to read the programs. Examples of computer readable media are hard drives, NAS systems, read-only memory (ROM), RAM, compact disks (CDs), digital versatile disks (DVDs), magnetic tapes, and other optical and non-optical data storage devices. A computer readable medium can also be distributed over a network-coupled computer system so that the computer readable code is stored and executed in a distributed fashion.

(52) Although one or more embodiments of the present invention have been described in some detail for clarity of understanding, certain changes may be made within the scope of the claims. Accordingly, the described embodiments are to be considered as illustrative and not restrictive, and the scope of the claims is not to be limited to details given herein but may be modified within the scope and equivalents of the claims. In the claims, elements and/or steps do not imply any particular order of operation unless explicitly stated in the claims.

(53) Virtualization systems in accordance with the various embodiments may be implemented as hosted embodiments, non-hosted embodiments, or as embodiments that blur distinctions between the two. Furthermore, various virtualization operations may be wholly or partially implemented in hardware. For example, a hardware implementation may employ a look-up table for modification of storage access requests to secure non-disk data.

(54) Many variations, additions, and improvements are possible, regardless of the degree of virtualization. The virtualization software can therefore include components of a host, console, or guest OS that perform virtualization functions.

(55) Plural instances may be provided for components, operations, or structures described herein as a single instance. Boundaries between components, operations, and data stores are somewhat arbitrary, and particular operations are illustrated in the context of specific illustrative configurations. Other allocations of functionality are envisioned and may fall within the scope of the invention. In general, structures and functionalities presented as separate components in exemplary configurations may be implemented as a combined structure or component. Similarly, structures and functionalities presented as a single component may be implemented as separate components. These and other variations, additions, and improvements may fall within the scope of the appended claims.

Claims

1. A method of deploying persistent storage for containers executing in a host cluster and managed by an orchestration control plane, the method comprising: receiving, at the orchestration control plane, a specification for implementation of an application in the containers and deployment of first persistent storage having a first storage class and second persistent storage having a second storage class; deploying, by the orchestration control plane, a first persistent volume as the first persistent storage in a virtual storage area network (SAN) as dictated by the first storage class, the virtual SAN comprising a plurality of local storage devices in the host cluster and implementing a policy that determines placement of the first persistent volume across the plurality of storage devices; and deploying, by the orchestration control plane, a second persistent volume as the second persistent storage in direct-attach storage as dictated by the second storage class, the direct-attach storage comprising other local storage devices in the host cluster not part of the virtual SAN, the specification including annotations that determine placement of the second persistent volume across the other local storage devices.

2. The method of claim 1, further comprising: deploying, by the orchestration control plane, the application executing on a first VM.

3. The method of claim 2, wherein the step of deploying the application comprises: specifying a replica count for the application; and specifying a persistent volume claim with a storage class, the storage class indicating of the second storage class, the second persistent volume being bound to the persistent volume claim.

4. The method of claim 3, wherein the step of deploying the application comprises: specifying an anti-affinity rule between the persistent volume claim and another persistent volume claim for the application.

5. The method of claim 1, further comprising: adding, by cooperation of a virtualization management server with the orchestration control plane, a persistent service for the orchestration

control plane, the application being an instance of the persistent service.

6. The method of claim 1, wherein a state of the orchestration control plane includes objects representing the host cluster and the first and second persistent volumes.

7. The method of claim 1, wherein the application stores the persistent data on the first and second persistent volumes.

8. A non-transitory computer readable medium comprising instructions to be executed in a computing device to cause the computing device to carry out a method of deploying persistent storage for containers executing in a host cluster and managed by an orchestration control plane, the method comprising: receiving, at the orchestration control plane, a specification for implementation of an application in the containers and deployment of first persistent storage having a first storage class and second persistent storage having a second storage class; deploying, by the orchestration control plane, a first persistent volume as the first persistent storage in a virtual storage area network (SAN) as dictated by the first storage class, the virtual SAN comprising a plurality of local storage devices in the host cluster and implementing a policy that determines placement of the first persistent volume across the plurality of storage devices; and deploying, by the orchestration control plane, a second persistent volume as the second persistent storage in direct-attach storage as dictated by the second storage class, the direct-attach storage comprising other local storage devices in the host cluster not part of the virtual SAN, the specification including annotations that determine placement of the second persistent volume across the other local storage devices.

9. The non-transitory computer readable medium of claim 8, further comprising: deploying, by the orchestration control plane, the application executing on a first VM.

10. The non-transitory computer readable medium of claim 9, wherein the step of deploying the application comprises: specifying a replica count for the application; and specifying a persistent volume claim with a storage class, the storage class indicating of the second storage class, the second persistent volume being bound to the persistent volume claim.

11. The non-transitory computer readable medium of claim 10, wherein the step of deploying the application comprises: specifying an anti-affinity rule between the persistent volume claim and another persistent volume claim for the application.

12. The non-transitory computer readable medium of claim 8, further comprising: adding, by cooperation of a virtualization management server with the orchestration control plane, a persistent service for the orchestration control plane, the application being an instance of the persistent service.

13. The non-transitory computer readable medium of claim 8, wherein a state of the orchestration control plane includes objects representing the host cluster and the first and second persistent volumes.

14. The non-transitory computer readable medium of claim 8, wherein the application stores the persistent data on the first and second persistent volumes.

15. A virtualized computing system, comprising: a cluster of hosts having a virtualization layer executing thereon and configured to manage virtual machines (VMs), the cluster of hosts comprising hardware platforms having random access memories (RAMs) and a plurality of local storage devices, the plurality of local storage devices configured as a virtual storage area network; another local storage device in a first host of the cluster, the other local storage device configured as direct-attach storage and not being one of the plurality of local storage devices; and an orchestration control plane, integrated with the virtualization layer, the orchestration control plane configured to: receive a specification for implementation of an application in containers and deployment of first persistent storage having a first storage class and second persistent storage having a second storage class; deploy a first persistent volume as the first persistent storage in the virtual storage area network as dictated by the first storage class; and deploy a second persistent volume as the second persistent storage in direct-attach storage as dictated by the second storage

class, the direct-attach storage comprising the other local storage device, the specification including annotations that determine placement of the second persistent volume on the other local storage device.

16. The virtualized computing system of claim 15, wherein the orchestration control plane is configured to: deploy the application executing on a first VM.

17. The virtualized computing system of claim 16, wherein the orchestration control plane is configured to: specify a replica count for the application; and specify a persistent volume claim with a storage class, the storage class indicating of the second storage class, the second persistent volume being bound to the persistent volume claim.

18. The virtualized computing system of claim 17, wherein the orchestration control plane is configured to: specify an anti-affinity rule between the persistent volume claim and another persistent volume claim for the application.

19. The virtualized computing system of claim 15, wherein the application stores the persistent data on the first and second persistent volumes.

20. The virtualized computing system of claim 15, wherein a state of the orchestration control plane includes objects representing the host cluster and the first and second persistent volumes.
