

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12405746
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Montero; Adolfo S. et al.

Local storage access by out-of-band (OOB) processors in heterogeneous computing platforms

Abstract

Systems and methods for local storage access by Out-of-Band (OOB) processors in heterogeneous computing platforms. In some embodiments, a heterogeneous computing platform may include an OOB Microcontroller Unit (MCU) and a memory coupled to the OOB MCU, the memory having program instructions stored thereon that, upon execution by the OOB MCU, cause the OOB MCU to access a local storage device external to the heterogeneous computing platform through an Embedded Controller (EC).

Inventors: Montero; Adolfo S. (Pflugerville, TX), Teshome; Abeye (Austin, TX), Pant; Alok (Austin, TX)

Applicant: Dell Products, L.P. (Round Rock, TX)

Family ID: 94387190

Assignee: Dell Products, L.P. (Round Rock, TX)

Appl. No.: 18/363879

Filed: August 02, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20250044978 A1	Feb. 06, 2025

Publication Classification

Int. Cl.: G06F3/06 (20060101)

U.S. Cl.:

Field of Classification Search

CPC: G06F (3/0604); G06F (3/0655); G06F (3/0679)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
11416028	12/2021	Kamepalli	N/A	G06F 1/3293
2010/0192145	12/2009	Liles	717/174	G06F 8/60
2016/0202964	12/2015	Butcher	717/172	G06F 8/654
2016/0232019	12/2015	Shah	N/A	G06F 9/45558
2017/0103225	12/2016	Pearson	N/A	G06F 21/79
2018/0324162	12/2017	Teshome	N/A	H04L 63/08
2021/0397577	12/2020	Kahn	N/A	H04L 67/1097

Primary Examiner: MacKall; Larry T

Attorney, Agent or Firm: Fogarty LLP

Background/Summary

FIELD

(1) This disclosure relates generally to Information Handling Systems (IHSs), and more specifically, to systems and methods for local storage access by Out-of-Band (OOB) processors in heterogeneous computing platforms.

BACKGROUND

(2) As the value and use of information continues to increase, individuals and businesses seek additional ways to process and store it. One option available to users is an Information Handling System (IHS). An IHS generally processes, compiles, stores, and/or communicates information or data for business, personal, or other purposes thereby allowing users to take advantage of the value of the information. Because technology and information handling needs and requirements vary between different users or applications, IHSs may also vary regarding what information is handled, how the information is handled, how much information is processed, stored, or communicated, and how quickly and efficiently the information may be processed, stored, or communicated.

(3) Variations in IHSs allow for IHSs to be general or configured for a specific user or specific use, such as financial transaction processing, airline reservations, enterprise data storage, or global communications. In addition, IHSs may include a variety of hardware and software components that may be configured to process, store, and communicate information and may include one or more computer systems, data storage systems, and networking systems.

(4) Historically, IHSs with desktop and laptop form factors have had conventional Operating Systems (OSs) (e.g., WINDOWS, LINUX, MAC OS, etc.) executed on INTEL or AMD's "x86"-type processors. Other types of processors, such as ARM processors, have been used in smartphones and tablet devices, which typically run thinner, simpler, or mobile OSs (e.g., ANDROID, IOS, WINDOWS MOBILE, etc.).

(5) More recently, however, IHS manufacturers have started producing full-fledged desktop and laptop IHSs equipped with ARM-based platforms. In fact, certain OSs (e.g., WINDOWS on ARM) have also been developed to provide users with a more quintessential OS experience on those platforms.

(6) Devices known as Embedded Controllers (ECs) have played a central role in their overall operation of traditional x86-based platforms. An EC is a microcontroller or processing core mounted on an IHS's motherboard which is configured to manage several critical IHS processes, ranging from early power rail sequencing to power limits and thermal limits, and to provide low-level hardware controls via a myriad of General-Purpose Input/Outputs (GPIOs).

(7) The EC is also responsible for facilitating Out-of-Band (OOB) management of its IHS. OOB management involves the use of dedicated interfaces for accessing and managing aspects of an IHS from a remote location, through a plane separate from the production network. As the inventors hereof have recognized, however, ARM-based platforms currently lack ECs (or other microcontrollers) usable to support OOB management.

SUMMARY

(8) Systems and methods for local storage access by Out-of-Band (OOB) processors in heterogeneous computing platforms are described. In an illustrative, non-limiting embodiment, a heterogeneous computing platform may include an OOB Microcontroller Unit (MCU) and a memory coupled to the OOB MCU, the memory having program instructions stored thereon that, upon execution by the OOB MCU, cause the OOB MCU to access a local storage device external to the heterogeneous computing platform through an Embedded Controller (EC).

(9) For example, the heterogeneous computing platform may include: a System-On-Chip (SoC), a Field-Programmable Gate Array (FPGA), or an Application-Specific Integrated Circuit (ASIC). The heterogeneous computing platform may include a Reduced Instruction Set Computer (RISC) processor and a plurality of devices coupled to an interconnect. The plurality of devices may include at least one of: a Graphical Processing Unit (GPU), an audio Digital Signal Processor (aDSP), a sensor hub, a Neural Processing Unit (NPU), a Tensor Processing Unit (TPU), a Neural Network Processor (NNP), an Intelligence Processing Unit (IPU), an Image Signal Processor (ISP), or a Video Processing Unit (VPU). And the interconnect may include at least one of: an Advanced Microcontroller Bus Architecture (AMBA) bus, a QuickPath Interconnect (QPI) bus, or a HyperTransport (HT) bus.

(10) The local storage device may include at least one of a Non-Volatile Memory Express (NVMe) device or a Solid-State Drive (SSD). The EC may be coupled to the local storage device via a sideband interface. The sideband interface may include an Inter-Integrated Circuit (I^{sup}.2C) bus or an Improved Inter-Integrated Circuit (I^{sup}.3C) bus. The OOB MCU may be configured to, while a host processor of the heterogeneous computing platform is in a low-power state: receive an OOB command, and in response to an indication that the OOB command requires a read or write operation, wake up the local storage device from a low-power state through the EC.

(11) The indication may include an indication that the read or write command involves at least one of: a Non-Volatile Memory Express (NVMe) boot partition, an NVMe protected additional partition, an Operating System (OS) boot file, an OS kernel file, an OS application file, a user file, or a security file. The OOB MCU may be further configured to access the local storage device through a high-bandwidth bus controller integrated into the heterogeneous computing platform.

(12) In another illustrative, non-limiting embodiment, an IHS may include a storage device and an OOB MCU integrated into a heterogeneous computing platform and coupled to the storage device through a high-bandwidth bus controller. The storage device may include an NVMe device or an SSD. The high-bandwidth bus controller may include a Peripheral Component Interconnect Express (PCIe) controller.

(13) The OOB MCU may be further configured to access the local storage device through an EC. To access the local storage device through the EC, the OOB MCU may be coupled to the EC

through a low-bandwidth bus controller. The low-bandwidth bus controller may include an I.sup.2C bus or an I.sup.3C bus.

(14) In another illustrative, non-limiting embodiment, a method may include: receiving, at a heterogeneous computing platform of an IHS, an OOB packet including a read or write command with respect to a local storage device coupled to the heterogeneous computing platform; and processing, by an EC coupled to the heterogeneous computing platform, the read or write command while a host processor of the heterogeneous computing platform is in a low-power state.

(15) The OOB packet may be retrieved by a network device coupled to the heterogeneous computing platform in response to an OOB packet sniffing operation while the host processor is in the low-power state. The read or write command may include an instruction to reimage a host OS executable by the host processor.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The present invention(s) is/are illustrated by way of example and is/are not limited by the accompanying figures, in which like references indicate similar elements. Elements in the figures are illustrated for simplicity and clarity, and have not necessarily been drawn to scale.

(2) FIG. 1 is a diagram illustrating examples of components of an Information Handling System (IHS), according to some embodiments.

(3) FIG. 2 is a diagram illustrating an example of a heterogenous computing platform, according to some embodiments.

(4) FIG. 3 is a diagram illustrating an example of a host Operating System (OS) executable by a heterogenous computing platform, according to some embodiments.

(5) FIG. 4 is a diagram illustrating an example of host OS-independent, autonomous inter-device communications in a heterogenous computing platform, according to some embodiments.

(6) FIGS. 5A-C are diagrams illustrating examples of different degrees of integration of an Embedded Controller (EC) in a heterogenous computing platform, according to some embodiments.

(7) FIG. 6 is a diagram illustrating an example of a heterogenous computing platform architecture for enabling Out-of-Band (OOB) management, according to some embodiments.

(8) FIG. 7 is a diagram illustrating an example of a heterogenous computing platform architecture for EC-based OOB packet sniffing from a low-power state, according to some embodiments.

(9) FIG. 8 is a diagram illustrating an example of a method for EC-based OOB packet sniffing from a low-power state, according to some embodiments.

(10) FIG. 9 is a diagram illustrating an example of a heterogenous computing platform architecture for networking-based OOB packet sniffing from a low-power state, according to some embodiments.

(11) FIG. 10 is a diagram illustrating an example of a method for networking-based OOB packet sniffing from a low-power state, according to some embodiments.

(12) FIG. 11 is a diagram illustrating OOB packet sniffing operations, according to some embodiments.

(13) FIG. 12 is a diagram illustrating an example of a system for caching OOB manageability packets, commands, or messages, according to some embodiments.

(14) FIG. 13 is a diagram illustrating an example of a method for caching OOB manageability packets, commands, or messages, according to some embodiments.

(15) FIG. 14 is a diagram illustrating an example of an Information Technology Decision Maker (ITDM) terminal with baseline and opaque OOB packets, according to some embodiments.

(16) FIG. 15 is a diagram illustrating an example of an opaque OOB packet, according to some

embodiments.

(17) FIG. 16 is a diagram illustrating an example of a manageability application, according to some embodiments.

(18) FIG. 17 is a diagram illustrating an example of a method for handling normal and opaque OOB packets, according to some embodiments.

(19) FIG. 18 is a diagram illustrating an example of a method for routing opaque OOB packets, according to some embodiments.

(20) FIG. 19 is a diagram illustrating an example of a system for handling opaque OOB packets in a sandbox, according to some embodiments.

(21) FIG. 20 is a diagram illustrating an example of a method for handling opaque OOB packets in a sandbox, according to some embodiments.

(22) FIG. 21 is a diagram illustrating an example of a power domain architecture for a heterogenous computing platform configured to perform OOB packet sniffing operations, according to some embodiments.

(23) FIG. 22 is a diagram illustrating an example of a system for progressive power control in a heterogenous computing platform, according to some embodiments.

(24) FIG. 23 is a diagram illustrating an example of a method for progressive power control in a heterogenous computing platform, according to some embodiments.

(25) FIG. 24 is a diagram illustrating an example of an interconnect architecture suitable for OOB manageability in a heterogenous computing platform, according to some embodiments.

(26) FIG. 25 is a diagram illustrating an example of a method for using an interconnect architecture to perform manageability operations, according to some embodiments.

(27) FIG. 26 is a diagram illustrating an example of a system for local storage access by OOB processors in a heterogenous computing platform, according to some embodiments.

(28) FIG. 27 is a diagram illustrating an example of a method for local storage access by OOB processors in heterogenous computing platforms, according to some embodiments.

DETAILED DESCRIPTION

(29) For purposes of this disclosure, an Information Handling System (IHS) may include any instrumentality or aggregate of instrumentalities operable to compute, calculate, determine, classify, process, transmit, receive, retrieve, originate, switch, store, display, communicate, manifest, detect, record, reproduce, handle, or utilize any form of information, intelligence, or data for business, scientific, control, or other purposes. For example, an IHS may be a personal computer (e.g., desktop or laptop), tablet computer, mobile device (e.g., Personal Digital Assistant (PDA) or smart phone), server (e.g., blade server or rack server), a network storage device, or any other suitable device and may vary in size, shape, performance, functionality, and price.

(30) An IHS may include Random Access Memory (RAM), one or more processing resources such as a Central Processing Unit (CPU) or hardware or software control logic, Read-Only Memory (ROM), and/or other types of nonvolatile memory. Additional components of an IHS may include one or more disk drives, one or more network ports for communicating with external devices as well as various I/O devices, such as a keyboard, a mouse, touchscreen, and/or a video display. An IHS may also include one or more buses operable to transmit communications between the various hardware components.

(31) The terms “heterogenous computing platform,” “heterogenous processor,” or “heterogenous platform,” as used herein, refer to an Integrated Circuit (IC) or chip (e.g., a System-On-Chip or “SoC,” a Field-Programmable Gate Array or “FPGA,” an Application-Specific Integrated Circuit or “ASIC,” etc.) containing a plurality of discrete processing circuits or semiconductor Intellectual Property (IP) cores (collectively referred to as “SoC devices” or simply “devices”) in a single electronic or semiconductor package, where each device has different processing capabilities suitable for handling a specific type of computational task. Examples of heterogenous processors include, but are not limited to: QUALCOMM's SNAPDRAGON, SAMSUNG's EXYNOS,

APPLE's "A" SERIES, etc.

(32) The term "firmware," as used herein, refers to a class of program instructions that provides low-level control for a device's hardware. Generally, firmware enables basic functions of a device and/or provides hardware abstraction services to higher-level software, such as an Operating System (OS). The term "firmware installation package," as used herein, refers to program instructions that, upon execution, deploy device drivers or services in an IHS or IHS component.

(33) The term "device driver" or "driver," as used herein, refers to program instructions that operate or control a particular type of device. A driver provides a software interface to hardware devices, enabling an OS and other applications to access hardware functions without needing to know precise details about the hardware being used. When an application invokes a routine in a driver, the driver issues commands to a corresponding device. Once the device sends data back to the driver, the driver may invoke certain routines in the application. Generally, device drivers are hardware dependent and OS-specific.

(34) The term "telemetry data," as used herein, refers to information resulting from in situ collection of measurements or other data by devices within a heterogeneous computing platform, or any other IHS device or component, and its transmission (e.g., automatically) to a receiving entity, for example, for monitoring purposes. Typically, telemetry data may include, but is not limited to, measurements, metrics, and/or values which may be indicative of: core utilization, memory utilization, network quality and utilization, battery utilization, peripheral or I/O device utilization, etc.

(35) For instance, telemetry data may include, but is not limited to, measurements, metrics, logs, or other information related to: current or average utilization of IHS components, CPU/core loads, instant or average power consumption of IHS components, instant or average memory usage by IHS components, characteristics of a network or radio system (e.g., WiFi vs. 5G, bandwidth, latency, etc.), keyboard or mouse usage data, transaction times, latencies, response codes, errors, data produced by other sensors, etc.

(36) To its user, an IHS may appear to be either "on" or "off," without any other detectable power states. In many implementations, however, an IHS may support multiple power states that correspond to the states defined in the Advanced Configuration and Power Interface (ACPI) specification, such as: S0, S1, S2, S3, S4, S5, and G3.

(37) For example, when an IHS is operating in S0 mode ("working"), the IHS is fully usable, yet hardware components not in use can still save power by entering a lower power state. In the S0 low-power idle mode ("Sleep" or "Modern Standby") an IHS remains partially running so that it can stay up to date whenever a suitable network is available, wake into S0 mode when real-time action is required (e.g., OS maintenance, etc.), and so on. In this state, an IHS can quickly switch from to a higher-power state in response to hardware and network events.

(38) In the past, S3 used to be the default "Sleep state;" however, now most IHSs have transitioned to "modern standby," which designates "S0ix" as a hybrid mode of S0, such that all devices are put into their lowest power state but x86 can still execute code such as networking for "connected standby."

(39) When an IHS does not support modern standby, it can instead assume any of states S1-S3 (also referred to as "Sleep" modes), where to a user the IHS may appear to be off. Some IHSs support only one or two of these Sleep states. Particularly, the amount of power consumed in states S1-S3 is less than S0 and more than S4. S3 consumes less power than S2, and S2 consumes less power than S1. In states S1-S3, volatile memory is kept refreshed to maintain the IHSs state, and some components remain powered so the IHS can wake from input from the keyboard, Local Area Network (LAN), or a Universal Serial Bus (USB) device. "Hybrid sleep," used only on certain types of IHSs, takes place when the IHS uses a hibernation file with S1-S3. The hibernation file saves the IHS's state in case it loses power while in sleep.

(40) In the S4 state ("Hibernate"), power consumption is reduced to its lowest level. The IHS saves

the contents of volatile memory to a hibernation file and some components remain powered so the computer can wake from input from the keyboard, LAN, or a USB device. The working context can be restored if stored on nonvolatile media. “Fast startup” is where the user is logged off before the hibernation file is created, which allows for a smaller hibernation file in IHSs with reduced storage capabilities.

(41) When in the S5 state (“Soft off” or “Full Shutdown”), an IHS is fully shut down without a hibernation file. It occurs when a restart is requested or when an application calls a shutdown Application Programming Interface (API) command. During a full shutdown and boot, the entire user session is torn down and restarted on the next boot, and a boot/startup from S5 takes significantly longer than S1-S4. At the hardware level, the only difference between S4 and S5 is that S4 keeps a flag on the storage device telling the OS bootloader to boot from the hibernation file instead of booting from scratch.

(42) In G3 (“Mechanical off”), the IHS is completely turned off and consumes absolutely no power from its Power Supply Unit (PSU) or main battery (e.g., a lithium-ion battery), with the exception of any Real-Time Clock (RTC) batter(ies) (e.g., Complementary Metal Oxide Semiconductor or “CMOS” batteries, Basic Input/Output System or “BIOS” batteries, coin cell batteries, etc.), which are used to provide power for the IHS's internal clock/calendar and for maintaining certain configuration settings. In some cases, G3 represents the lowest possible power configuration of an IHS from which the IHS can be initialized; all voltage rails are turned off except for RTC rails.

(43) From G3 an IHS typically goes to S5 based on AC insertion (for transition between battery mode to AC mode). Additionally, or alternatively, an IHS may transition from G3 to S0 based upon the detection of a power button event.

(44) As used herein, the term “Out-of-Band (OOB) management” refers to the use of dedicated interfaces for accessing and managing aspects (e.g., firmware updates, change BIOS settings, collect device inventory, collect telemetry data, etc.) of an IHS from a remote location, through a plane separate from the production network. For example, a form of out-of-band management is referred to as Lights-Out Management (LOM), which employs a dedicated management channel for device maintenance and allows a system administrator to monitor and manage an IHS by remote control, regardless of whether the IHS is powered on or whether an OS is installed or functional.

(45) It should be noted that the term “OOB management” is used in contradistinction with “in-band management,” the latter being based upon in-band connectivity (the usual or “production” network channel) and requiring that software be installed on the IHS being managed, which works only after the OS has been successfully booted and networking is brought up.

(46) FIG. 1 is a block diagram of components of IHS 100. As depicted, IHS 100 includes host processor(s) 101. In various embodiments, IHS 100 may be a single-processor system, or a multi-processor system including two or more processors. Host processor(s) 101 may include any processor capable of executing program instructions, such as an INTEL/AMD x86 processor, or any general-purpose or embedded processor implementing any of a variety of Instruction Set Architectures (ISAs), such as a Complex Instruction Set Computer (CISC) ISA, a Reduced Instruction Set Computer (RISC) ISA (e.g., one or more ARM core(s), or the like).

(47) IHS 100 includes chipset 102 coupled to host processor(s) 101. Chipset 102 may provide host processor(s) 101 with access to several resources. In some cases, chipset 102 may utilize a QuickPath Interconnect (QPI) bus to communicate with host processor(s) 101. Chipset 102 may also be coupled to communication interface(s) 105 to enable communications between IHS 100 and various wired and/or wireless networks, such as Ethernet, WiFi, BT, cellular or mobile networks (e.g., Code-Division Multiple Access or “CDMA,” Time-Division Multiple Access or “TDMA,” Long-Term Evolution or “LTE,” etc.), satellite networks, or the like.

(48) Communication interface(s) 105 may be used to communicate with peripherals devices (e.g., BT speakers, microphones, headsets, etc.). Moreover, communication interface(s) 105 may be

coupled to chipset **102** via a Peripheral Component Interconnect Express (PCIe) bus, or the like.

(49) Chipset **102** may be coupled to display and/or touchscreen controller(s) **104**, which may include one or more Graphics Processor Units (GPUs) on a graphics bus, such as an Accelerated Graphics Port (AGP) or PCIe bus. As shown, display controller(s) **104** provide video or display signals to one or more display device(s) **111**.

(50) Display device(s) **111** may include Liquid Crystal Display (LCD), Light Emitting Diode (LED), organic LED (OLED), or other thin film display technologies. Display device(s) **111** may include a plurality of pixels arranged in a matrix, configured to display visual information, such as text, two-dimensional images, video, three-dimensional images, etc. In some cases, display device(s) **111** may be provided as a single continuous display, rather than two discrete displays.

(51) Chipset **102** may provide host processor(s) **101** and/or display controller(s) **104** with access to system memory **103**. In various embodiments, system memory **103** may be implemented using any suitable memory technology, such as static RAM (SRAM), dynamic RAM (DRAM) or magnetic disks, or any nonvolatile/Flash-type memory, such as a Solid-State Drive (SSD), Non-Volatile Memory Express (NVMe), or the like.

(52) In certain embodiments, chipset **102** may also provide host processor(s) **101** with access to one or more USB ports **108**, to which one or more peripheral devices may be coupled (e.g., integrated or external webcams, microphones, speakers, etc.).

(53) Chipset **102** may further provide host processor(s) **101** with access to one or more hard disk drives, solid-state drives, optical drives, or other removable-media drives **113**.

(54) Chipset **102** may also provide access to one or more user input devices **106**, for example, using a super I/O controller or the like. Examples of user input devices **106** include, but are not limited to, microphone(s) **114A**, camera(s) **114B**, and keyboard/mouse **114N**. Other user input devices **106** may include a touchpad, stylus or active pen, totem, etc. Each of user input devices **106** may include a respective controller (e.g., a touchpad may have its own touchpad controller) that interfaces with chipset **102** through a wired or wireless connection (e.g., via communication interfaces(s) **105**). In some cases, chipset **102** may also provide access to one or more user output devices (e.g., video projectors, paper printers, 3D printers, loudspeakers, audio headsets, Virtual/Augmented Reality (VR/AR) devices, etc.).

(55) In certain embodiments, chipset **102** may further provide an interface for communications with one or more hardware sensors **110**. Sensors **110** may be disposed on or within the chassis of IHS **100**, or otherwise coupled to IHS **100**, and may include, but are not limited to: electric, magnetic, radio, optical (e.g., camera, webcam, etc.), infrared, thermal, force, pressure, acoustic (e.g., microphone), ultrasonic, proximity, position, deformation, bending, direction, movement, velocity, rotation, gyroscope, Inertial Measurement Unit (IMU), and/or acceleration sensor(s).

(56) BIOS **107** is coupled to chipset **102**. UEFI was designed as a successor to BIOS, and many modern IHSs utilize UEFI in addition to or instead of a BIOS. Accordingly, BIOS/UEFI **107** is intended to also encompass a UEFI component BIOS/UEFI **107** provides an abstraction layer that allows the OS to interface with certain hardware components that are utilized by IHS **100**.

(57) Upon booting of IHS **100**, host processor(s) **101** may utilize program instructions of BIOS **107** to initialize and test hardware components coupled to IHS **100**, and to load host OS **300** for use by IHS **100**. Via the hardware abstraction layer provided by BIOS/UEFI **107**, software stored in system memory **103** and executed by host processor(s) **101** can interface with certain I/O devices that are coupled to IHS **100**.

(58) Embedded Controller (EC) **109** (sometimes referred to as a Baseboard Management Controller or “BMC”) includes a microcontroller unit or processing core dedicated to handling selected IHS operations not ordinarily handled by host processor(s) **101**. Examples of such operations may include, but are not limited to: power sequencing, power management, receiving and processing signals from a keyboard or touchpad, as well as other buttons and switches (e.g., power button, laptop lid switch, etc.), receiving and processing thermal measurements (e.g., performing cooling

fan control, CPU and GPU throttling, and emergency shutdown), controlling indicator Light-Emitting Diodes or “LEDs” (e.g., caps lock, scroll lock, num lock, battery, ac, power, wireless LAN, sleep, etc.), managing the battery charger and the battery, enabling remote management, diagnostics, and remediation over network(s) **103**, etc.

(59) Unlike other devices in IHS **100**, EC **109** may be made operational from the very start of each power reset, before other devices are fully running or powered on. As such, EC **109** may be responsible for interfacing with a power adapter to manage the power consumption of IHS **100**. These operations may be utilized to determine the power status of IHS **100**, such as whether IHS **100** is operating from battery power or is plugged into an AC power source. Firmware instructions utilized by EC **109** may be used to manage other core operations of IHS **100** (e.g., turbo modes, maximum operating clock frequencies of certain components, etc.).

(60) In some cases, EC **109** may implement operations for detecting certain changes to the physical configuration or posture of IHS **100** and managing other devices in different configurations of IHS **100**. For instance, when IHS **100** as a 2-in-1 laptop/tablet form factor, EC **109** may receive inputs from a lid position or hinge angle sensor **110**, and it may use those inputs to determine: whether the two sides of IHS **100** have been latched together to a closed position or a tablet position, the magnitude of a hinge or lid angle, etc. In response to these changes, the EC may enable or disable certain features of IHS **100** (e.g., front or rear facing camera, etc.).

(61) In this manner, EC **109** may identify any number of IHS postures, including, but not limited to: laptop, stand, tablet, or book. For example, when display(s) **111** of IHS **100** is open with respect to a horizontal keyboard portion, and the keyboard is facing up, EC **109** may determine IHS **100** to be in a laptop posture. When display(s) **111** of IHS **100** is open with respect to the horizontal keyboard portion, but the keyboard is facing down (e.g., its keys are against the top surface of a table), EC **109** may determine IHS **100** to be in a stand posture. When the back of display(s) **111** is closed against the back of the keyboard portion, EC **109** may determine IHS **100** to be in a tablet posture. When IHS **100** has two display(s) **111** open side-by-side, EC **109** may determine IHS **100** to be in a book posture. In some implementations, EC **109** may also determine if display(s) **111** of IHS **100** are in a landscape or portrait orientation.

(62) In some implementations, EC **109** may be installed as a Trusted Execution Environment (TEE) component to the motherboard of IHS **100**. Additionally, or alternatively, EC **109** may be further configured to calculate hashes or signatures that uniquely identify individual components of IHS **100**. In such scenarios, EC **109** may calculate a hash value based on the configuration of a hardware and/or software component coupled to IHS **100**. For instance, EC **109** may calculate a hash value based on all firmware and other code or settings stored in an onboard memory of a hardware component.

(63) Hash values may be calculated as part of a trusted process of manufacturing IHS **100** and may be maintained in secure storage as a reference signature. EC **109** may later recalculate the hash value for a component may compare it against the reference hash value to determine if any modifications have been made to the component, thus indicating that the component has been compromised. As such, EC **109** may validate the integrity of hardware and software components installed in IHS **100**.

(64) In addition, EC **109** may provide an OOB channel that allows an Information Technology Decision Maker (ITDM) or Original Equipment Manufacturer (OEM) to manage IHS **100**'s various settings and configurations, for example, by issuing OOB commands.

(65) In various embodiments, IHS **100** may be coupled to an external power source through an AC adapter, power brick, or the like. The AC adapter may be removably coupled to a battery charge controller to provide IHS **100** with a source of DC power provided by battery cells of a battery system in the form of a battery pack (e.g., a lithium ion or “Li-ion” battery pack, or a nickel metal hydride or “NiMH” battery pack including one or more rechargeable batteries).

(66) Battery Management Unit (BMU) **112** may be coupled to EC **109** and it may include, for

example, an Analog Front End (AFE), storage (e.g., non-volatile memory), and a microcontroller. In some cases, BMU **112** may be configured to collect and store information, and to provide that information to other IHS components, such as, for example devices within heterogeneous computing platform **200** (FIG. 2).

(67) Examples of information collectible by BMU **112** may include, but are not limited to: operating conditions (e.g., battery operating conditions including battery state information such as battery current amplitude and/or current direction, battery voltage, battery charge cycles, battery state of charge, battery state of health, battery temperature, battery usage data such as charging and discharging data; and/or IHS operating conditions such as processor operating speed data, system power management and cooling system settings, state of “system present” pin signal), environmental or contextual information (e.g., such as ambient temperature, relative humidity, system geolocation measured by GPS or triangulation, time and date, etc.), events, etc.

(68) Examples of events may include, but are not limited to: acceleration or shock events, system transportation events, exposure to elevated temperature for extended time periods, high discharge current rate, combinations of battery voltage, battery current and/or battery temperature (e.g., elevated temperature event at full charge and/or high voltage causes more battery degradation than lower voltage), etc.

(69) In some embodiments, IHS **100** may not include all the components shown in FIG. 1. In other embodiments, IHS **100** may include other components in addition to those that are shown in FIG. 1. Furthermore, some components that are represented as separate components in FIG. 1 may instead be integrated with other components, such that all or a portion of the operations executed by the illustrated components may instead be executed by the integrated component.

(70) For example, in various embodiments described herein, host processor(s) **101** and/or other components shown in FIG. 1 (e.g., chipset **102**, display controller(s) **104**, communication interface(s) **105**, EC **109**, etc.) may be replaced by devices within heterogeneous computing platform **200** (FIG. 2). As such, IHS **100** may assume different form factors including, but not limited to: servers, workstations, desktops, laptops, appliances, video game consoles, tablets, smartphones, etc.

(71) FIG. 2 is a diagram illustrating an example of heterogeneous computing platform **200**. In various embodiments, heterogeneous computing platform **200** may be implemented in an SoC, FPGA, ASIC, or the like. Heterogeneous computing platform **200** includes a plurality of discrete or segregated devices, each device having a different set of processing capabilities suitable for handling a particular type of computational task. When each device in platform **200** executes only the types of computational tasks it is specifically designed to execute, the overall power consumption of heterogeneous computing platform **200** is reduced.

(72) In various implementations, each device in heterogeneous computing platform **200** may include its own microcontroller(s) or core(s) (e.g., ARM core(s)) and corresponding firmware. In some cases, a device in platform **200** may also include its own hardware-embedded accelerator (e.g., a secondary or co-processing core coupled to a main core). Each device in heterogeneous computing platform **200** may be accessible through a respective Application Programming Interface (API). Additionally, or alternatively, each device in heterogeneous computing platform **200** may execute its own OS. Additionally, or alternatively, one or more of these devices may be a virtual device.

(73) In certain embodiments, at least one device in heterogeneous computing platform **200** may have updatable firmware which, upon installation, operates to change the performance, available features, settings, configuration options, API, drivers, and/or services provided by that device. For example, each update may be delivered to heterogeneous computing platform **200** as a system-wide firmware installation package having a plurality of firmware components, and each firmware component may be distributed to its respective device (or corresponding memory space). In some implementations, the latest system-wide firmware installation package received by heterogeneous computing platform **200** may be installed at every boot of IHS **100**.

(74) In the example of FIG. 2, heterogenous computing platform **200** includes CPU clusters **201A-N** as a particular implementation of host processor(s) **101** intended to perform general-purpose computing operations. Each of CPU clusters **201A-N** may include one or more processing core(s) and cache memor(ies). In operation, CPU clusters **201A-N** are available and accessible to the IHS's host OS **300** (e.g., WINDOWS on ARM), optimization application(s) **312** (FIG. 3), OS agent(s) **313**, and other application(s) **314** executed by IHS **100**.

(75) CPU clusters **201A-N** are coupled to memory controller **202** via internal interconnect fabric **203**. Memory controller **202** is responsible for managing memory accesses for all of devices connected to internal interconnect fabric **203**, which may include any communication bus suitable for inter-device communications within an SoC (e.g., Advanced Microcontroller Bus Architecture or "AMBA," QuickPath Interconnect or "QPI," HyperTransport or "HT," etc.). All devices coupled to internal interconnect fabric **203** can communicate with each other and with a host OS executed by CPU clusters **201A-N**.

(76) GPU **204** is a device designed to produce graphical or visual content and to communicate that content to a monitor or display, where the content may be rendered. PCIe interfaces **205** provide an entry point into any additional devices external to heterogenous computing platform **200** that have a respective PCIe interface (e.g., graphics cards, USB controllers, etc.).

(77) Audio Digital Signal Processor (aDSP) **206** is a device designed to perform audio and speech operations and to perform in-line enhancements for audio input(s) and output(s). Examples of audio and speech operations include, but are not limited to: noise reduction, echo cancellation, directional audio detection, wake word detection, muting and volume controls, filters and effects, etc.

(78) In operation, input and/or output audio streams may pass through and be processed by aDSP **206**, which can send the processed audio to other devices on internal interconnect fabric **203** (e.g., CPU clusters **201A-N**). Also, aDSP **206** may be configured to process one or more of heterogenous computing platform **200**'s sensor signals (e.g., gyroscope, accelerometer, pressure, temperature, etc.), low-power vision or camera streams (e.g., for user presence detection, onlooker detection, etc.), or battery data (e.g., to calculate a charge or discharge rate, current charge level, etc.). To that end, aDSP **206** may be coupled to BMU **112**.

(79) Sensor hub and low-power AI device **207** is a very low power, always-on device designed to consolidate information received from other devices in heterogenous computing platform **200**, process any context and/or telemetry data streams, and provide that information to: (i) a host OS, (ii) other applications, and/or (iii) other devices in platform **200**. For example, sensor hub and low-power AI device **207** may include General-Purpose Input/Output (GPIOs) that provide Inter-Integrated Circuit (I.sup.2C), Improved I.sup.2C (I.sup.3C), Serial Peripheral Interface (SPI), Enhanced SPI (eSPI), and/or serial interfaces to receive data from sensors (e.g., sensors **110**, camera **210**, peripherals **214**, etc.).

(80) Sensor hub and low-power AI device **207** may include an always-on, low-power core configured to execute small neural networks and specific applications, such as contextual awareness and other enhancements. In some embodiments, sensor hub and low-power AI device **207** may be configured to operate as an orchestrator device in charge of managing other devices, for example, based upon an ITDM/OEM policy or the like.

(81) High-performance AI device **208** is a significantly more powerful processing device than sensor hub and low-power AI device **207**, and it may be designed to execute multiple complex AI algorithms and models concurrently (e.g., Natural Language Processing, speech recognition, speech-to-text transcription, video processing, gesture recognition, user engagement determinations, etc.). For example, high-performance AI device **208** may include a Neural Processing Unit (NPU), Tensor Processing Unit (TPU), Neural Network Processor (NNP), or Intelligence Processing Unit (IPU), and it may be designed specifically for AI and Machine Learning (ML), which speeds up the processing of AI/ML tasks while also freeing processor(s) **101** to perform other tasks.

(82) Display/graphics device **209** is designed to perform additional video enhancement operations. In operation, display/graphics device **209** may provide a video signal to an external display coupled to IHS **100** (e.g., display device(s) **111**).

(83) Camera device **210** includes an Image Signal Processor (ISP) configured to receive and process video frames captured by a camera coupled to heterogeneous computing platform **200** (e.g., in the visible and/or infrared spectrum).

(84) Video Processing Unit (VPU) **211** is a device designed to perform hardware video encoding and decoding operations, thus accelerating the operation of camera **210** and display/graphics device **209**. VPU **211** may be configured to provide optimized communications with camera device **210** for performance improvements.

(85) In some cases, devices **209-211** may be coupled to internal interconnect fabric **203** via a secondary interconnect fabric (not shown). A secondary interconnect fabric may include any bus suitable for inter-device and/or inter-bus communications within an SoC.

(86) Security device **212** includes any suitable security device, such as a dedicated security processor, a Trusted Platform Module (TPM), a TRUSTZONE device, a PLUTON processor, or the like. In various implementations, security device **212** may be used to perform cryptography operations (e.g., generation of key pairs, validation of digital certificates, etc.) and/or it may serve as a hardware root-of-trust (RoT) for heterogeneous computing platform **200** and/or IHS **100**.

(87) Modem/wireless controller **213** is a device designed to enable all wireless communications in any suitable frequency band (e.g., BLUETOOTH or “BT,” WiFi, CDMA, 5G, satellite, etc.), subject to AI-powered optimizations/customizations for improved speeds, reliability, and/or coverage.

(88) Peripherals **214** may include any device coupled to heterogeneous computing platform **200** (e.g., sensors **110**) through mechanisms other than PCIe interfaces **205**. In some cases, peripherals **214** may include interfaces to integrated devices (e.g., built-in microphones, speakers, and/or cameras), wired devices (e.g., external microphones, speakers, and/or cameras, Head-Mounted Devices/Displays or “HMDs,” printers, displays, etc.), and/or wireless devices (e.g., wireless audio headsets, etc.) coupled to IHS **100**.

(89) In some cases, devices **212** and **213** may be coupled to internal interconnect fabric **203** via the same secondary interconnect serving devices **209-211** (not shown). Additionally, or alternatively, devices **212-213** may be coupled to internal interconnect fabric **203** via another secondary interconnect.

(90) In various embodiments, one or more devices of heterogeneous computing platform **200** (e.g., GPU **204**, aDSP **206**, sensor hub and low-power AI device **207**, high-performance AI device **208**, VPU **211**, etc.) may be configured to execute one or more AI model(s), simulation(s), and/or inference(s).

(91) FIG. 3 is a diagram illustrating an example of host OS **300** executable by CPU clusters or cores **201A-N** of heterogeneous computing platform **200**. In some cases, host OS **300** may be implemented as WINDOWS on ARM. As shown, host OS **300**'s stack includes kernel mode drivers (KMD) in kernel space **301** and user mode drivers (UMD) in user space **302**.

(92) In kernel space **301**, OS secure kernel **303** is responsible for secure operations (e.g., encryption, validation, etc.) within IHS **100**. Core OS/API service **304** has direct access to IHS components such as, for example, heterogeneous computing platform **200**. OS drivers **305** include kernel mode drivers developed by the OS publisher or other developer. Platform drivers **306** include kernel mode drivers developed by the manufacturer or vendor of heterogeneous computing platform **200**, for example, for use with devices **201-214**.

(93) In user space **302**, user-mode platform drivers and services **307** enable access to features provided by devices **201-214** through kernel-mode platform drivers **306**. OEM drivers **308** enable features in OEM devices coupled to IHS **100**, and user-mode OS drivers and services **309** enable access to OS features through kernel mode OS drivers **305**. Platform secure kernel **310** includes

protected user-mode portions of host OS **300** developed by the manufacturer or vendor of heterogenous computing platform **200**, and OS secure kernel extensions **311** include extensions to protected user-mode portions of host OS **300** developed by the OS publisher or other developer.

(94) OS agent(s) **313** may include an OS agent or client configured to communicate with service(s) provided by an ITDM/OEM service to establish an in-band management channel. Moreover, other application(s) **314** may include or be a part of any workload executable by heterogenous computing platform **200**. For example, other application(s) **314** may include productivity, collaboration, streaming, multimedia, or gaming applications executable by host OS **300**.

(95) Optimization application(s) **312** may include any host OS **300**-level application configured to intelligently optimize the performance of IHS **100** (e.g., DELL OPTIMIZER), for example, by using AI models to dynamically configure IHS **100**'s settings and/or to optimize the performance of other applications **314**. In operation, optimization application(s) **312** may improve the productivity, performance, and user experience through system usage analysis and learning. For example, optimization application(s) **312** may be invoked, within host OS **300**, to learn how application **314** is used. Optimization application(s) **312** may identify optimization opportunities, classify users, and automatically apply appropriate settings (e.g., storage, memory, and/or CPU) to improve application **314**'s performance.

(96) At least one of applications **312-314** may be configured to utilize one or more devices, features, or services exposed, surfaced, enumerated, or otherwise made available by user-mode drivers **307-309**, for example, through a Human Interface Device (HID) interface and/or an HID report provided by host OS **300**, or the like.

(97) In various implementations, host OS **300** may be configured to receive a firmware installation package or executable file containing at least one extension driver (e.g., a setup information or "INF" text file in a driver package) from an ITDM/OEM service via a management channel. The installation package may be processed by a UEFI UpdateCapsule process when distributed as part of an OS update, as a system-wide and/or device-specific firmware update, and/or it may be distributed by OEM update applications such as DELL COMMAND UPDATE, integrated with remote deployment and update management tools.

(98) Firmware components of each extension driver may be loaded, attached, or extended onto user-mode platform drivers and services **307**, and may be communicated by user-mode platform drivers and services **307** to respective devices of heterogenous computing platform **200** through kernel-mode platform drivers **306** for installation, update, or execution of such firmware components in those devices.

(99) As such, the deployment of extension drivers by host OS **300** as OEM drivers **308** provides value-added features or services to user-mode platform drivers (e.g., base drivers) **307** and/or applications **312-314**. Particularly, OEM drivers **308** may expose custom services and routines provided by any one of devices **201-214** upon execution of their corresponding firmware components.

(100) FIG. 4 is a diagram illustrating an example of host OS-independent, autonomous inter-device communications **400** in heterogenous computing platform **200**. Particularly, each of devices **401A-N** may implement any of devices **201-214** of heterogenous computing platform **200**.

(101) Each of APIs **402A-N** provides access to firmware **403A-N** executed by a corresponding device **401A-N**. In operation, each firmware component **403A-N** may exchange data and commands directly with each other using APIs **402A-N**. Through APIs **402A-N**, one or more of devices **401A-N** may be configured to enable, disable, or modify firmware services provided by other devices **403A-N**. For example, in some embodiments, one or more of devices **401A-N** may be designated as an orchestrator upon booting of platform **200**.

(102) In various embodiments, firmware services resulting from the execution of firmware **403A-N** may be provided by their respective device **401A-N** to other devices **401A-N** within heterogeneous computing platform **200** autonomously and/or independently of the operation of host OS **300**.

Communications between executing firmware **403A-N** and applications **312-314** may take place through OEM drivers **308**. In some cases, executing firmware **403A-N** may be identified by or exposed to host OS **300** and/or applications **312-314** as part of an HID report.

(103) Firmware **403A-N** and corresponding OEM drivers **308** may be installed, modified, updated, and/or removed from IHS **100** upon each installation of a firmware installation package for the entire heterogenous computing platform **200**, for instance, at each boot of IHS **100**. For example, each firmware component **403A-N** providing a respective firmware service may be delivered to a respective device **401A-N** as an extension driver. Upon execution, these firmware services may provide additional controls over the management, deployment, customization, and/or configuration of IHS **100** to an ITDM/OEM that are independent of updates to host OS **300** and/or applications **312-314**.

(104) In other embodiments, any given one of devices **401A-N** may be rebooted or reset independently of the other devices to perform a local installation, update, or modification of that device's firmware services without having to reboot the entire heterogenous computing platform **200** and/or IHS **100**. Additionally, or alternatively, one or more of devices **401A-N** may have its firmware service at least partially installed or updated without rebooting or resetting the device.

(105) As discussed above, systems and methods described herein may integrate EC **109** into heterogenous computing platform **200** of IHS **100**. For example, in the implementations shown in FIGS. 5A and 5B, components of EC **109** (e.g., core, peripheral devices, etc.) may be fully or partially integrated into heterogeneous computing platform **200**. In another implementation shown in FIG. 5C, however, EC **109** may be completely external to platform **200** but coupled to a bridge integrated therein via an interface (e.g., enhanced SPI or "eSPI") to provide or maintain the EC's ability to access the SoC's internal interconnect fabric. In each of these scenarios, EC **109** may be configured to operate as an orchestrator instead of (or along with) sensor hub and low-power AI device **207**.

(106) Particularly, FIG. 5A is a diagram illustrating example **500A** of EC **109A** integrated into heterogenous computing platform **200A**. In this embodiment, EC **109A** is fully integrated within the SoC and it is coupled to internal interconnect fabric **203** via internal bus or interconnect access **501**. EC **109A** may include EC processing core **502A** (e.g., another ARM core) and a plurality of peripheral devices **503A-N**.

(107) Examples of peripheral devices **503A-N** include, but are not limited to: an Inter-Integrated Circuit (I2C) bus controller, a Pulse Width Modulator (PWM), a tachometer (TACH), a GPIO controller, a serial peripheral interface (SPI) Flash, a fan controller, an Analog-to-Digital Converter (ADC), etc.

(108) Each of peripheral devices **503A-N** may be coupled to a corresponding pin or terminal of SoC **200A**. The operating voltage of heterogenous computing platform **200A** is typically lower than the operating voltage of external EC-managed devices (e.g., cooling fans, etc.), which means that, ordinarily, peripheral devices **503A-N** would not be able to communicate with them. To address this problem, voltage translator **505** may be coupled to peripheral devices **503A-N** through low-voltage connections **504A-N** (via corresponding pins of the SoC) and to EC-managed devices (not shown) via high-voltage connections **506A-N**.

(109) In operation, voltage translator **505** may be configured to increase the voltage of signals transmitted from peripheral devices **503A-N** to EC-managed devices (e.g., from 1.2 or 1.8 V to 3.3 V). Voltage translator **505** may also be configured to reduce the voltage of other signals transmitted from EC-managed devices to peripheral devices **503A-N** (e.g., from 3.3 V to 1.8 or 1.2 V).

(110) FIG. 5B shows a diagram of example **500B** of EC **109B** partially integrated into heterogenous computing platform **200B**. In contrast with example **500A**, here only EC core **502B** is integrated into the SoC and peripheral devices **508A-N** (e.g., similar to peripheral devices **503A-N**, but with different operating voltages) are provided in separate package **509** mounted on the IHS motherboard, apart from the SoC.

(111) To communicate with peripheral devices **503A-N**, EC core **502B** may be coupled to package **509** via connection or bus **507**, such as a SPI bus (e.g., 1.2 V, 100 MHz). In this manner, example **500B** reduces the number of pins (e.g., ~**100** fewer pins or more) on the SoC that would otherwise be required to enable all EC operations. Moreover, because peripheral devices **503A-N** are external to EC **109B**, they may interface with EC-managed devices directly over **506A-N** using their native voltages (e.g., 3.3 V), and therefore voltage translator **505** of example **500A** is not required.

(112) As such, in examples **500A** and **500B**, EC **109A/B** may power up as soon as SoC **200A/B** comes out of reset. In some cases, Microcontroller Unit (MCU) **500A/B** may own a local SRAM program store that is loaded once when coming out of reset, and which does not rely on an external DRAM to be initialized and able to run. In other implementations, such internal memory may be Flash-based.

(113) Accordingly, EC **109A/B** may have access to all the SoC **200A/B**'s internal resources, and especially early access to the SoC **200A/B**'s GPIOs for early hardware bring-up control points, such as power rails or other activities otherwise typically reserved for a Power Management Controller (PMIC). Moreover, EC **109A/B** may have access to the SoC **200A/B**'s internal hardware interconnect bus **203** to enable special or enhanced operations not currently possible with x86 architectures.

(114) FIG. 5C shows a diagram of example **500C** of EC **109C** externally disposed with respect to heterogenous computing platform **200C**. In contrast with examples **500A/B**, here both EC core **502C** and peripheral devices **512A-N** are provided in a discrete package mounted on the IHS motherboard. The SoC includes bridge **511** coupled to internal interconnect fabric **203** via internal bus or interconnect access **510**, and EC **109C** is coupled to bridge **511** via high-speed bus or interface **512**, which may be an eSPI interface or the like (such as a derivative of PCIe type protocols at the hardware level).

(115) In example **500C**, the deployment of eSPI bus **512** may prevent bottlenecks in the EC's operations that would result if other types of interfaces (e.g., I^{sup}.2C) were used. For instance, eSPI bus **512** between the SoC and EC **109C** may enable the EC **109C** to run multiple parallel features concurrently. Also, eSPI bus **512** provide the capability to virtualize GPIOs using "virtual wires," whereby EC **109C** may access the GPIO status over eSPI instead of going through the physical GPIO hardware layer.

(116) Furthermore, high-speed interface **512** may provide EC **109C** with full connectivity to the internal SoC fabric. For example, in a scenario where host processor(s) **101** fails to boot, eSPI bus **512** may allow EC **109C** to access and/or run most or all of devices **401A-N** (e.g., PCIe bus, USB controllers, network controllers, graphics controllers, system memory, etc.) directly.

(117) As such, example **500C** may be used, for example, to provide pre-boot, on-demand services based upon EC **109C**; which is not currently possible with existing SoCs (or in x86 architectures). EC **109C** is generally verifiable to be a trusted device, thus once that trust is established, having access to key SoC resources no longer presents attack vector concerns. For example, in various implementations, EC **109C** may have sideband read/write access to hard drive **113**, an SSD, or an NVMe coupled to platform **200**.

(118) A possible benefit or advantage of enabling external EC **109C** over internal ECs **109A/B** is the ability to operate GPIOs at different voltage levels than the SoC. Another possible benefit or advantage of external EC **109C** versus internal ECs **109A/B** includes reducing the GPIO pin count of the SoC, especially in implementations where the SoC's footprint is already saturated with GPIOs and a designer is forced to grow the size of the SoC to accommodate potentially hundreds of additional GPIO pinouts.

(119) In various embodiments, systems and methods described herein may enable OOB management of heterogeneous computing platform **200** in low-power states (e.g., S5, G3, any power state where only the RTC voltage rails are on, etc.).

(120) In that regard, FIG. 6 is a diagram illustrating an example of heterogenous computing

platform architecture **600** for enabling OOB management. In this example, heterogeneous computing platform **200** of FIG. 2 or the like may be implemented in the form of SoC **601** within IHS **100**, such that host processor(s) **101** are implemented as host processor(s) **602** having main cores (or application processors) **603A-N** configured to execute BIOS/UEFI **107** and host OS **300**. (121) SoC **601** includes integrated OOB MCU **604**. OOB MCU **604** may include a processing core integrated into SoC **601** but that is otherwise separate and independent from any of main cores **603A-N** of processor **602**. In various embodiments, OOB MCU **604** may belong to its own power domain, and it may be more power efficient than main cores **603A-N** (typically with less processing capabilities), but less power efficient than EC **109** (typically with greater processing capabilities). In contrast with host processor(s) **602**, OOB MCU **604** may not run any components of host OS **300** (or BIOS/UEFI **107**), and it may be invisible to any host OS **300** and/or or BIOS/UEFI **107** running in main cores **603A-N**.

(122) As shown, OOB MCU **604** may be coupled to (or include): OOB boot Read-Only Memory (ROM) **605**, integrated static RAM (SRAM) **606**, memory controller **607**, and network traffic routing controller **610**.

(123) OOB ROM **605** includes program instructions and/or firmware that enable OOB MCU **604** to start and maintain its operations (e.g., Real-Time OS or “RTOS” code), including certain operations described in more detail below. Integrated SRAM **606** is usable by OOB MCU **604** as its working memory during such operations.

(124) Within IHS **100**, memory controller **607** may provide OOB MCU **604** with access to system memory **103**, here implemented as soldered down memory **608** (e.g., fixed Double Data Rate Synchronous Dynamic Random-Access Memory (ies) or “DDR SDRAM” memory(ies)) and/or Dual In-line Memory Modules or “DIMMs” **609** (e.g., removable DDR memory(ies)) within IHS **100** for additional processing capabilities.

(125) Network traffic routing controller **610** is configured to manage both in-band network traffic from host-processor cores **602** (e.g., generated or processed by host OS **300**) and OOB network traffic from OOB MCU **604** (e.g., generated or processed by executing firmware stored in OOB ROM **605**), for example, when IHS **100** and/or SoC **601** are both in S0, using any suitable arbitration technique.

(126) When operating in G3 or S5, the path between network traffic routing controller **610** and host processor(s) **602** is off, and network traffic routing controller **610** exclusively handles OOB traffic to/from OOB MCU **604**.

(127) In cases where BIOS/UEFI **107** is launched during a G3 Sniff operation and requires the use of network device **614**, however, network traffic routing controller **610** may arbitrate inbound and outbound traffic between OOB MCU **604** and host processor(s) **602** while SoC **601** is otherwise in G3 (or S5).

(128) To manage in-band and OOB packet traffic, packets may be interleaved using an interleaving policy. For example, such interleaving policy may include one or more rules dictating that network traffic routing controller **610** route a maximum number or amount of OOB packets or data to or from OOB MCU **604** only after routing a minimum number or amount of in-band packets or data to or from host processor(s) **602**.

(129) Moreover, in various implementations, in order to make a distinction between packets directed to OOB MCU **604** and/or host processor(s) **602**, network traffic routing controller **610** may be configured to examine the header of each packet and determine whether it is an OOB packet or an in-band packet based upon a header indication. The header indication may inform network traffic routing controller **610** that the packet is an OOB packet, and/or it may include destination information (e.g., an identification of OOB MCU **604**).

(130) In some cases, network traffic routing controller **610** may be configured to apply a routing priority policy to determine which packets should be routed to its respective destination first. For example, network traffic routing controller **610** may store a look-up table of OOB commands or

types of commands that, when included in an OOB packet, give that OOB packet routing priority over in-band packets (e.g., a “kill pill” command).

(131) Network traffic routing controller **610** provides OOB MCU **604** with access to low-bandwidth bus controller **611** and high-bandwidth bus controller **613**. In operation, network traffic routing controller **610** may be configured to arbitrate and/or multiplex in-band packet traffic to/from host processor(s) **602** with OOB packet traffic to/from OOB MCU **604**.

(132) In architecture **600**, OOB MCU **604** may be coupled to network traffic routing controller **610** via a first portion of a high-speed bus (e.g., PCIe, USB Type-C or 3.0, I.sup.3C, eSPI, etc.), network traffic routing controller **610** may be coupled to high-bandwidth bus controller **613** via a second portion of the high-speed bus, and high-bandwidth bus controller **613** may be coupled to network device **614** via a third portion of the high-speed bus. In some cases, the first, second, and/or third portions of the high-speed bus may use the same or a different protocol.

(133) In some cases, the use of a high-speed bus or protocol for incoming and outgoing communications between OOB MCU **604** and network device **614** (independent of the power state of host processor(s) **602**, such as during a G3 Sniff operation) may prevent bottlenecks in the OOB MCU **604**'s operations that would result if other types of interfaces (e.g., slow-speed interfaces, such as I.sup.2C) were used.

(134) In various embodiments, network traffic routing controller **610** may provide OOB MCU **604** with the ability to share the same wireless device (e.g., network device **614**) that IHS **100** uses for host OS **300**'s runtime. When in a low-power modes or pre-boot mode, OOB MCU **604** fully owns the wireless/wired interface; however, in cases where host OS **300** is up and running, port addressing logic within network traffic routing controller **610** may redirect traffic dedicated to OOB MCU **604** for receiving packets, and OOB MCU **604** may also be able to send outgoing packets to be merged with host OS **300**-level network traffic, even when host OS **300** is running, while protecting the OOB network traffic from OS sniffing at the hardware level.

(135) In situations where host OS **300** hangs (e.g., Blue Screen of Death or “BSOD”), OOB MCU **604** may take over network traffic routing controller **610**. A dual-ported memory may allow simultaneous access to the network device **614** by both network traffic routing controller **610** and host OS **300**'s drivers.

(136) Moreover, power state transitions may alert or notify OOB MCU **604**, upon or preceding entry into Modern Standby, that OOB MCU **604** may make low power use of the network layer without waking host processor(s) **602** and/or host OS **300** (e.g., less OOB network traffic during Modern Standby than during S0).

(137) Low-bandwidth bus controller **611** (e.g., I.sup.2C, I.sup.3C, Low-Pin Count or “LPC,” etc.) is coupled to flash memory **612**, which is configured to store EC firmware, BIOS firmware, etc., and to EC **109**. EC **109** may also be coupled directly to flash memory **612** through any suitable interface.

(138) High-bandwidth bus controller **613** (e.g., PCIe interface **205**, USB Type-C or 3.0, etc.) is coupled to networking hardware, component(s), or device(s) **614** (e.g., modem, wireless controller, LAN, WiFi, cellular, satellite, BT, etc.). Network device **614** may be capable of communicating with cloud server(s) **617** via wired connection **615** (e.g., Ethernet, fiber optics, etc.) and/or wireless connection **616** (e.g., WiFi, cellular, etc.). EC **109** may also have a direct connection to network device **614** via sideband bus **620** (e.g., I.sup.3C, eSPI, etc.).

(139) In various embodiments, the terms low-bandwidth bus controller **611** and high-bandwidth bus controller **613** are meant to illustrate that the data speeds of controller **613** are greater than the bus speeds of controller **611**; sometimes by two or more orders of magnitude (e.g., PCIe versus I.sup.2C).

(140) Cloud server(s) **617** may be accessible by ITDM terminal or console **618**, which executes or otherwise includes manageability application **619**. In some cases, cloud server(s) **617** and ITDM terminal **618** may each be implemented as different instances of IHS **100**.

(141) In some implementations, EC **109** may be at least partially integrated into SoC **601** (e.g., FIGS. 5A and 5B). In other implementations, flash memory **612** and network device **614** may also be integrated as one of devices within SoC **601**. In yet other implementations, OOB MCU **604** may be disposed outside of SoC **601** (e.g., within EC **109** or on its own, as a discrete processor or controller).

(142) The inventors hereof have identified a desire to support OOB manageability capabilities, particularly when IHS **100** and/or SoC **601** is in a low-power state, such as S5 or G3. The G3 state in x86 designs for mobile IHSs is currently defined to as when AC is absent and the IHS is turned off. Desktop IHSs predominantly remain in S5 mode because its AC power source is generally continuously present, except in special cases where the IHS is disposed in a location prone to blackouts, or in applications where all Point-of-Sale (POS) IHSs are switched off on a nightly basis, for example. Currently, in the G3 state, an EC does not run any firmware and the IHS is generally unable to support any OOB manageability operations.

(143) To address these issues, systems and methods described herein enable EC **109** to perform “G3 Sniffing,” “Low-power Sniffing,” “OOB Sniffing,” and/or “OOB packet Sniffing” operations, whereby EC **109** comes up periodically out of G3 state using an RTC timer to “sniff” for any OOB manageability packets, commands, or messages targeted to that individual IHS and stored in cloud server(s) **617**.

(144) The duty cycle of EC **109**'s sniff operations may be selected to be anywhere from a very low duty cycle (e.g., less than 2% of the time) to a very high duty cycle (e.g., more than 70% of the time), with a tradeoff between responsiveness and power consumption.

(145) Unlike other technologies where an MCU wakes up following a given duty cycle in order to reduce a device's power consumption, systems and methods described herein enable EC **109** to temporarily wake up from ultra-low power mode (e.g., G3) for the purpose of sniffing OOB manageability packets, commands, or messages, and to return to such ultra-low power mode upon completion of the G3 sniff.

(146) Using these systems and methods, EC **109** may employ an RTC timer to wake periodically (e.g., at the expiration or completion of the timer's countdown) to check, using network device **614**, for any pending OOB manageability packets, commands, or messages addressed to that particular IHS. Moreover, power rail controls may be implemented such that during G3 sniff state, when EC comes out of G3 via RTC timer, the power island or domain of network device **614** is also energized to allow it to check cloud service(s) **617** for any pending OOB manageability packets, commands, or messages.

(147) FIG. 7 is a diagram illustrating an example of heterogeneous computing platform architecture **700** for EC-based OOB packet sniffing from a low-power state. In this embodiment, EC **109** is coupled to RTC power **704** (e.g., a CMOS battery) and 3V_ALW power rail **703** (an “always-on” voltage rail), and it includes RTC timer **701** and RTC wake logic **702**.

(148) EC **109** is also coupled to power control Field-Effect Transistor (FET) or switch **705**, which in turn is coupled to network device **614**.

(149) While in G3 mode, EC **109** is configured to receive power from RTC power **704**. When RTC timer **701** is triggered, it sends a signal to RTC wake logic **702** which wakes up EC **109**, at least in part, by switching its power from RTC power **704** to 3V_ALW power **703**. When EC **109** is sufficiently operational, it sends a GPIO signal to power control FET **705**, which also brings network device **614** out of G3 so that EC **109** may trigger and/or facilitate, through sideband bus **620**, network device **614**'s sniff operations from cloud service(s) **617** through wired connection **615** or wireless connection **616**.

(150) If there are no OOB manageability packets, commands, or messages to be retrieved and/or performed, or upon completion of these operations, EC **109** and network device **614** may return to G3. Other devices within SoC **601** not involved during such operations (e.g., host processor(s) **602**) may remain in G3 mode during the G3 sniffing.

(151) Although shown as being performed by EC **109**, in other situations, the entity performing the aforementioned operations may be OOB MCU **604**, in which case OOB MCU **604** may include blocks **701-704**, and it may be coupled to power control FET **705** to wake up network device **614**.

(152) FIG. **8** is a diagram illustrating an example of method **800** for EC-based OOB packet sniffing from a low-power state. In various embodiments, method **800** may be performed, at least in part, by EC **109** in cooperation with network device **614**.

(153) Particularly, method **800** starts at **801**. At **802**, IHS **100** is in a low power state, such as G3, receiving power from RTC power **704**. At **803**, RTC timer **701** issues a signal to RTC wake logic **702**, which proceeds to switch EC **109** from RTC power **704** to 3V_ALW power **703**.

(154) At **804**, EC **109** issues a GPIO signal to power control FET **705** to wake up network device **614**. In some cases, at **804**, the same (or a different) GPIO signal may also wake up network traffic routing controller **610**.

(155) At **805**, EC **109** may send and receive messages to/from cloud server(s) **617** via wired connection **615** or wireless connection **616**, including querying for OOB manageability packets, commands, or messages targeting its IHS, for example, by checking the status of a command buffer (e.g., by sending a message asking whether such buffer exists and/or how many OOB manageability packets, commands, or messages have been queued) associated with that IHS.

(156) At **806**, EC **109** determines whether cloud server(s) **617** has OOB manageability packets, commands, or messages queued or buffered for its IHS (e.g., by receiving a response or acknowledgement from cloud server(s) **617**). If not, at **807**, EC **109** turns off power control FET **705** and both EC **109** and network traffic routing controller **610** return to G3. EC **109** may restart its RTC timer **701**, disconnect from 3V_ALW power **703**, and connect to RTC power **704**.

(157) If so, at **808**, EC **109** may process OOB manageability packets, commands, or messages retrieved from cloud server(s) **617** before returning to G3 at **809**. In some cases, such a packet may request execution of a command that requires EC **109** to wake up other components in SoC **601** such as, for example, OOB MCU **604** (e.g., via another power control FET).

(158) Additionally, or alternatively, if a command involves a BIOS or OS update, for example, EC **109** may wake up host processor(s) **602** (e.g., with another power control FET) and/or the entire SoC **601**. After **807** or **808**, control returns to **802**.

(159) An alternative approach to architecture **700** of FIG. **7** includes configuring network device **614** to wake itself up intermittently and to check for OOB manageability packets, commands, or messages targeted for IHS **100**. In those cases, power rail support allows both EC **109** and networking hardware **104** to remain in G3 enabled by the RTC rail until a periodic timer wakes networking hardware **104** and until networking hardware **104** wakes up EC **109**. Network hardware **614** may be disposed on a power island similar to the RTC low power rail and it may use an RTC low power timer to periodically wake from the G3 state (i.e., instead of EC **109** doing the wake from G3 from RTC power rail).

(160) After network device **614** wakes up, it may either: (1) go back to sleep if no OOB manageability packets, commands, or messages are pending for that IHS, or (2) send a wake alarm to EC **109** to indicate there is work to be done. The wake alarm to EC **109** may be communicated via a VBAT-Powered Control Interface (VCI) or GPIO pin/logic such that EC **109** may wake from the G3 state in order to process the pending management command in conjunction with network device **614**.

(161) FIG. **9** is a diagram illustrating an example of heterogenous computing platform architecture **900** for networking-based OOB packet sniffing from a low-power state. In this embodiment, network device **614** is coupled to RTC power **704** and 3V_ALW power rail **703** (an “always-on” voltage rail), and it includes RTC timer **901** and RTC wake logic **902**. Network device **614** is also coupled to VCI pin wake logic **930** within EC **109**.

(162) Network device **614** is configured to receive power from RTC power **704** while in G3 mode. When RTC timer **901** is triggered, it sends a signal to RTC wake logic **902** which wakes up

network device **614**, at least in part, by switching its power from RTC power **704** to 3V_ALW power **703**. When network device **614** is sufficiently operational, it may send a VCI signal to VCI pin wake logic **903**, which brings EC **109** out of G3 so that EC **109** may process packets received through network device **614**'s sniff operations from cloud service(s) **617** through wired connection **615** or wireless connection **616**.

(163) In various embodiments, network device **614** may be configured to only wake EC **109** if there are pending data packets found in cloud service **617**. If no packets associated with IHS **100** are pending, only network device **614** wakes up momentarily, then goes back to G3, without waking up EC **109**.

(164) Upon completion of OOB operations, EC **109** and network device **614** return to G3. Moreover, other devices within SoC **601** not required during such operations (e.g., host processor(s) **602**) may remain in G3 mode.

(165) In the case of wired networking (LAN), architecture **900** may also receive a “magic packet” and wake EC **109** to process the packet if the packet is addressed to EC **109**. Circuitry within network device **614** may be segmented to a very low power island only looking for magic packets, and additional power may be provided when the magic packet match has been confirmed by such circuitry.

(166) FIG. **10** is a diagram illustrating an example of method **1000** for networking-based OOB packet sniffing from a low-power state. In various embodiments, method **1000** may be performed, at least in part, by network device **614** in cooperation with EC **109**.

(167) Particularly, method **1000** starts at **1001**. At **1002**, IHS **100** is in a low power state, such as G3, receiving power from RTC power **704**. At **1003**, RTC timer **901** issues a signal to RTC wake logic **902**, which proceeds to switch network device **614** from RTC power **704** to 3V_ALW power **703**.

(168) At **1004**, network device **614** may send and receive messages to/from cloud server(s) **617** via wired connection **615** or wireless connection **616**, including sniffing for OOB manageability packets, commands, or messages, for example, by checking the status of a command buffer associated with its IHS.

(169) At **1005**, if OOB manageability packets, commands, or messages are not present, network device **614** may disconnect from 3V_ALW power **703** and connect to RTC power **704** to back to G3 state at **1006**, and it may restart RTC timer **901**.

(170) Conversely, if OOB manageability packets, commands, or messages are present at **1005**, network device **614** may send a VCI signal to VCI pin logic **903** to wake EC **109** out of G3 at **1007** so that EC **109** may process the received OOB manageability packets, commands, or messages at **1008**. Upon processing the OOB manageability packets, commands, or messages, EC **109** may also return to G3 at **1009**. After **1006** or **1009**, control returns to **1002**.

(171) A yet alternative approach to architectures **700** and **900** includes configuring OOB MCU **604** to wake itself up intermittently and to check for OOB manageability packets, commands, or messages targeted for IHS **100**, similar to EC **109**.

(172) FIG. **11** is a diagram illustrating OOB packet sniffing operations **1100**. In some embodiments, operations **1100** may be performed during the execution of method **800** and/or **1000**. Particularly, IHS **100** is in G3 mode at **1101A**, when EC **109**, OOB MCU **604**, or network device **614** wake up to perform sniff operation **1102A**. In this example, the result of sniff operation **1102A** (e.g., 250 ms duration) is that there are no OOB manageability packets, commands, or messages buffered for IHS **100** in cloud service(s) **617**; therefore EC **109**, OOB MCU **604**, or network device **614** return to G3 at **1101B**.

(173) EC **109**, OOB MCU **604**, or network device **614** again wake up to perform sniff operation **1102B** at the expiration of their RTC timer **701/901** (e.g., after 5 minutes, 10 minutes, 30 minutes, etc.) In this example, the result of sniff operation **1102B** is that there are OOB manageability packets, commands, or messages buffered for IHS **100** in cloud service(s) **617**; for instance, EC

109 may wake up other devices within SoC **601** at **1103** (e.g., OOB MCU **604**) to process the OOB manageability packets, commands, or messages at **1104**. At **1101C**, all components return to G3. (174) At a subsequent expiration of RTC timer **701/901**, EC **109**, OOB MCU **604**, or network device **614** wake up to perform sniff operation **1102C**. Again, the result of sniff operation **1102C** is that there are no OOB manageability packets, commands, or messages buffered for IHS **100** in cloud service(s) **617**; therefore EC **109** and/or network device **614** return to G3 at **1101D**. (175) The time interval between sniff operations **1102A-C** is selectable through configuration of RTC timer **701/901**. In some cases, such an interval may allow SoC **601** to remain in G3 state 90% of the time, or any other selected amount of time. (176) On the cloud server(s) **617** side, systems and methods described herein may enable the caching of OOB commands targeted to any particular IHS, such that multiple managed IHS have individual queues of what operations they need to carry out. In cases of immediate commands when OOB channel is continuously up, each command queue may be emptied out instantly. When IHS **100** is in G3, however, the cloud server(s) **617** may hold queues of IHS-specific commands and it may allow each IHS coming out of that state to perform a sniff operation configured to pull any pending manageability commands targeting IHS **100**. (177) FIG. **12** is a diagram illustrating an example of system **1200** for caching OOB manageability packets, commands, or messages. In various embodiments, system **1200** may be used to buffer OOB manageability packets, commands, or messages (“CMD1,” “CMD2,” etc.) addressed to a selected one of a plurality of IHSs while the IHS is in a low-power state (e.g., G3), before it performs a sniffing operation. Particularly, cloud server(s) **617** is coupled to one or more ITDM terminal(s) **618** configured to execute manageability application **619**. (178) Cloud server(s) **617** receive OOB manageability packets, commands, or messages from ITDM terminal **618** and stores those packets, commands, or messages in appropriate ones of distinct buffers **1201A-N**, each buffer associated with one of a plurality of IHSs **1202A-N**. Each of IHSs **1202A-N** (e.g., **100**) may include its own EC **1202A-N** (e.g., **109**) and G3 sniff network components **1204A-N** (e.g., network device **614**). (179) Moreover, each of buffers **1201A-N** may store a plurality of OOB manageability packets, commands, or messages (“CMD1” through “CMDN”). In some cases, each OOB manageability packet, command, or message in each buffer may be associated with a priority value (e.g., “0” or “1”, between “1” and “5,” etc.) that indicates an order in which that packet, command, or message should be retrieved by its respective IHS **1202A-N**, and the packets, commands, or messages in each buffer may be sorted by priority. In other cases, each of buffers **1201A-N** may provide stored OOB manageability packets, commands, or messages to its respective IHS following a first-in-first-out (FIFO) rule or technique, last-in-first-out (LIFO) rule or technique, etc. (180) Examples of OOB manageability packets, commands, or messages may include, but are not limited to: an IHS provisioning command, a locate IHS command, put the IHS in a selected power mode, restrict access to the IHS, entitlements, firmware updates, change BIOS settings, system or device inventory collection command, collect telemetry data, a kill pill, etc. (181) FIG. **13** is a diagram illustrating an example of method **1300** for caching OOB manageability packets, commands, or messages. In some embodiments, method **1300** may be performed by cloud server(s) **617** in cooperation with IHSs **1202A-N**. Particularly, method **1300** starts at **1301**. At **1302**, cloud server(s) **617** may receive an OOB manageability packet, command, or message from ITDM terminal **618** targeting a selected one of IHSs **1202A-N** (e.g., IHS “X”). At **1303**, cloud server(s) **617** determines whether IHS “X” is awake (e.g., by sending a message and waiting for an acknowledgement). (182) At **1303**, if cloud server(s) **617** determines that IHS “X” is awake (e.g., the acknowledgement is received), then at **1304** it transmits the OOB manageability packet, command, or message to IHS “X.” In some cases, if IHS “X” is determined to be awake (e.g., S0), cloud server(s) **617** may send a message to IHS “X” to prevent it from entering a sleep state until its OOB manageability packets,

commands, or messages have been processed.

(183) Conversely, at **1303**, if cloud server(s) **617** determines that IHS “X” is not awake (e.g., the acknowledgement is not received), then at **1305** it stores the OOB manageability packet, command, or message in a buffer **1201A-N** associated with IHS “X.” When IHS “X” partially wakes up to perform a G3 Sniff operation, it may follow methods **800** and/or **1000** to retrieve its queued OOB manageability packets, commands, or messages. After executing or processing all of its packets, commands, or messages, all components of IHS “X” may return to the G3 state.

(184) In that regard, the inventors hereof also have recognized that, when OOB packets are sent to a target IHS, there is a need to allow a separation between: (a) “normal,” “standard,” or “baseline” OOB packets, which contain management commands and features that a maker of SoC **601** may choose to enable for all IHS manufacturers (e.g., Original Equipment Manufacturer or “OEM”) without discrimination; and (b) “enhanced,” “special,” or “opaque” OOB packets, which contain management commands and features that are proprietary or otherwise specific to a given IHS OEM, in order to allow the OEM to provide differentiating features to its customers.

(185) To address these, and other concerns, systems and methods described herein may enable distinguishing baseline OOB packets from opaque OOB packets, such that the processing of baseline OOB packets may be covered by a publicly available specification. Conversely, opaque OOB packets may be processed by OEM-owned firmware without the maker of SoC **601** having any knowledge of at least a portion of its contents (e.g., the payload portion of the opaque OOB packet). In some embodiments, OEM-owned firmware may reside primarily within an EC **109**’s domain, or in a protected customizable space (an isolated execution area or “sandbox”) within OOB MCU **604**.

(186) FIG. **14** is a diagram illustrating example **1400** of ITDM terminal **618** with baseline OOB packets **1401** and opaque OOB packets **1402** handled by manageability application **619**. As such, example **1400** splits OOB traffic into separate two categories or types: a “baseline” type for public domain functionality available to all OEMs; and an “opaque” type for OEM-differentiated or special capabilities beyond the baseline—which may be treated as “opaque data” whereby the OOB infrastructure simply passes the packets through to OEM-owned firmware stacks.

(187) In other embodiments, three or more categories or types of OOB commands or packets **1402** may be used. For example, in addition to a baseline category for all OEMs, ITDM terminal **618** may enable the creation of two or more opaque packet categories: a first category or sub-type of opaque OOB packets for a specific OEM that relates to free services or features provided by that OEM, and a second category or sub-type of opaque OOB packets for the same OEM that relates to paid or premium features provided by that OEM.

(188) In some embodiments, different types of opaque OOB packets **1402** may be associated with a specific map, key, or decode ring that structurally indicates, bit-by-bit and/or field-by-field, what variable or setting each bit value in the opaque OOB packet is mapped to (that is, what the value means) for a given OEM. Moreover, the map, key, or decode ring necessary for interpreting opaque OOB packet **1402** may not be shared with the maker of SoC **601** (except, in some cases, for the header portion of opaque packet **1402**). Instead, such map, key, or decode ring may reside only within the OEM firmware (e.g., EC **109**), so that to any entity other than EC **109**, opaque OOB packet **1402** may remain undecipherable.

(189) FIG. **15** is a diagram illustrating an example of opaque OOB packet **1500**. In some embodiments, opaque packet **1500** may include opaque identifier **1501**, opaque routing header **1502**, and opaque payload **1503**. For example, opaque identifier **1501** may indicate whether packet **1500** is a baseline packet or an opaque packet (or a sub-type of opaque packet). Meanwhile, opaque routing header **1502** may indicate a target device, within a particular IHS (e.g., EC **109**, BIOS/UEFI **107**, OOB MCU **604**, host OS **300**, any device within SoC **601**, etc.) for processing opaque OOB packet **1500**.

(190) FIG. **16** is a diagram illustrating an example of manageability application **619**. In various

embodiments, manageability application **619** may include User Interface (UX) layer **1601** having baseline UX handling plugin **1602** and enhanced UX plug-ins **1603A-N**. Command processing module **1604** includes baseline command handler(s) **1605** and enhanced command plugin(s) **1606A-N**. Manageability application **619** also includes command I/O layer **1607** and networking layer **1608**, such that the latter is in communication with cloud server(s) **617**.

(191) In operation, an ITDM may access or execute manageability application **619**, for example, using ITDM terminal **618** coupled to cloud server(s) **617**. The ITDM may use baseline UX handling module **1602** or one of plugins **1603A-N** to enter one or more OOB management commands targeting a selected one of plurality of IHSs and/or a selected device or endpoint within the selected IHS. If the OOB management command is a baseline OOB command, module **1602** is invoked. Conversely, if the OOB management command is an opaque, OEM-specific command, it may be entered through one of plugins **1603A-N** (e.g., one plugin for each OEM, one plugin per type of opaque command, etc.).

(192) Once the OOB management command is received by user interface layer **1601**, it is passed to one of handlers **1605** or **1606A-N** within command processing module **1604**. In some cases, each of modules **1605** is associated with a corresponding one of plugin **1603A-N**.

(193) Each of handlers **1605** or **1606A-N** may be configured to take an OOB management command from a corresponding one of module **1602** or plugins **1603A-N** and to produce one or more baseline or opaque OOB packets containing the OOB management command, for example, using a unique packet map, key, or decode ring.

(194) Baseline OOB packets may have a packet map, key, or decode ring that is known, whereas opaque OOB packets may have packet maps, keys, or decode rings that are secret and known only to the OEM. Moreover, when a new OEM and/or opaque packet type is added, a respective plugin **1603A-N** and corresponding handler **1606A-N** may also be added to manageability application **619**.

(195) Handlers **1605** or **1606A-N** pass the OOB packets to command I/O layer **1607**. Command I/O layer **1607** forwards the OOB packets to networking layer **1608**, from where the OOB packets are sent to cloud server(s) **617** and then forwarded to a respective IHS, as identified in the packet header (e.g., opaque header **1502**), if the IHS is running. If the IHS is in a low-power mode (e.g., G3), however, cloud server(s) **617** may queue the OOB packet in an associated one of command buffers **1201A-N**.

(196) FIG. **17** is a diagram illustrating an example of method **1700** for handling normal and opaque OOB packets. In some embodiments, method **1700** may be performed, at least in part, by OOB MCU **604**. Particularly, method **1700** starts at **1701**.

(197) At **1702**, OOB MCU **604** may receive an OOB packet (from network device **614**) and it may inspect the OOB packet's opaque identifier **1501**, if any. If the OOB packet is not opaque, it may be processed by OOB MCU **604**. Conversely, if at **1702** method **1700** determines that the OOB packet is opaque, OOB MCU **604** may forward the OOB packet to EC **109**. After **1704** or **1705**, control returns to **1702**.

(198) In this example, EC **109** includes an opaque packet map, key, or decode ring for decoding or interpreting the opaque packet, and the EC's firmware under control of the IHS's OEM is the only entity with such a map, key, or decode ring. In other cases, however, OOB MCU **604** may include a sandbox portion (described below) configured to process opaque packets on behalf of the OEM without access by the manufacturer, maker, vendor, or supplier of SoC **601**, such that the sandbox portion of OOB MCU **614** includes the opaque packet map, key, or decode ring. The sandbox may constitute a separate or isolated execution area not visible to the SoC **601**'s vendor.

(199) Additionally, or alternatively, method **1700** may be performed by network device **614**, in which case network device **614** includes the opaque packet map, key, or decode ring.

(200) When EC **109** (or any other endpoint) needs to respond to the OOB management command included in an inbound opaque OOB packet, EC **109** may send such a response message using the

same opaque map, key, or decoder ring as the inbound OOB packet. The outbound response message may be at least partially undecipherable by OOB MCU **604** or any entity other than an OEM server configured to receive such a message. As such, the response message may leave EC **109** and reach network device **614** through OOB MCU **604** without its payload being decipherable. (201) Although described as being performed by OOB MCU **604**, in some cases method **1700** may instead be performed by network device **614**. For example, if SoC **601** is in G3 state, it may be more efficient to have network device **614** inspect the opaque OOB packet and, if it is intended to be executed by EC **109**, forward the opaque OOB packet to EC **109** without waking OOB MCU **604** up from G3.

(202) More generally, when OOB packets are sent to an IHS in the opaque data format, if all of them were routed only to EC **109**, many use cases would not be fully optimized because certain OOB management commands are performed by other entities (e.g., host processor(s) **602**), and therefore require additional data transfers from EC **109** to those entities.

(203) To address these, and other concerns, routing header **1502** of OOB packet **1500** may indicate a target or destination endpoint within SoC **601** of IHS **100**, such as, for example: EC **109**, BIOS/UEFI **107** (as executed by host processor(s) **602**), and OOB MCU **604**. OOB management operations that are best handled by EC **109** (and which are OEM-differentiated features would use the opaque data model and EC **109** as the target endpoint for the opaque data packet. OOB MCU **604** may ensure the data path chosen is the hardware link to EC **109** through a pre-established communications link between OOB MCU **604** and EC **109**.

(204) Similarly, OOB management commands that are best handled by the BIOS may be routed directly to the BIOS/UEFI **107** when host processor(s) **602** is up and running. OOB MCU **604** may use SoC data paths to send the packet to BIOS/UEFI **107** when BIOS/UEFI **107** is noted as the endpoint in routing header **1502**. Alternately if BIOS/UEFI **107** is not up and running, OOB MCU **604** may queue the OOB packet for later delivery, when BIOS/UEFI **107** may send a flag to OOB MCU **604** that BIOS/UEFI **107** is up and running. Finally, OOB management commands that are best handled by OOB MCU **604** itself, yet within the Dell proprietary context, may be routed to a sandbox or an OEM FW stack handler within OOB MCU **604**, as described in more detail below.

(205) FIG. **18** is a diagram illustrating an example of method **1800** for routing opaque OOB packets. In some embodiments, method **1800** may be performed, at least in part, by OOB MCU **604**. Particularly, method **1800** begins at **1801**.

(206) At **1802**, OOB MCU **604** may receive an OOB packet (e.g., from network device **614**). At **803**, OOB MCU **604** determines whether the OOB packet is baseline or opaque. If the OOB packet is a baseline packet, OOB MCU **604** processes the packet and executes the OOB management command contained therein at **1804**. Conversely, if the OOB packet is an opaque packet, as indicated in identifier **1501**, at **1805** method **1800** may check the contents of packet header **1502**.

(207) At **1806**, if OOB MCU **604** determines that the OOB packet's target is EC **109**, at **1810** OOB MCU **604** may forward the packet to EC **109** (e.g., via low-bandwidth bus controller **611**). At **1807**, if OOB MCU **604** determines that the OOB packet's target is BIOS/UEFI **107**, at **1811** OOB MCU **604** may forward the packet to BIOS/UEFI **107** (e.g., via host processor(s) **602**). At **1808**, if OOB MCU **604** determines that the OOB packet's target is OOB MCU **604**, at **1812** OOB MCU **604** may process the OOB packet using the OEM's "sandbox," as described below. Otherwise at **1809** method **1800** may invoke an error handling process.

(208) Although three OOB packet targets or destinations are specifically shown in method **1800**, in alternative embodiments any number of targets (e.g., host OS **300**, other devices within SoC **601**, etc.) may be specified in opaque header **1502**. Moreover, although method **1800** is shown as being performed by OOB MCU **604**, in other cases these routing operations may be performed by EC **109** and/or network device **614**.

(209) When an opaque OOB packet reaches OOB MCU **604**, OOB MCU **604** may need to separate: (a) a standard OOB MCU firmware, from (b) an OEM custom firmware, in order to

handle the opaque OOB packets separately from baseline OOB packets, so as not to compromise enhanced or exclusive manageability features added by the IHS's OEM.

(210) To address these, and other issues, systems and methods described herein may add a “sandbox” section to OOB MCU **604** such that OEM-specific firmware can have access to the hardware capabilities of OOB MCU **604**. The manufacturer or vendor of SoC **601** may provide a baseline API to their existing OOB MCU **604** capabilities and hardware peripherals, and an OEM can write any features or commands necessary for implementing its differentiating features (e.g., handling opaque packets routed to OOB MCU **604**) without disclosing those features to third parties.

(211) FIG. **19** is a diagram illustrating an example of system **1900** for handling opaque OOB packets in a sandbox. As shown, OOB MCU **604** may execute OOB baseline firmware **1901** and OEM sandbox custom firmware **1902** separately from each other.

(212) In some implementations, once OOB MCU **604** executes OOB baseline firmware **1901**, it may instantiate a main RTOS that accepts OEM sandbox custom firmware **1902** to operate as a segregated RTOS module. In other implementations, however, OEM sandbox custom firmware **1902** may be executed in other suitable ways such that it may have access to all resources and APIs that the manufacturer or vendor of SoC **601** provides as part of its baseline operations.

(213) As such, in some cases, OOB baseline firmware **1901** may be used to process baseline OOB commands or packets. Conversely, OEM sandbox custom firmware **1902** may be used to process opaque OOB commands or packets to the exclusion of OOB baseline firmware **1901**.

(214) To execute different firmware separately, OOB MCU **604** may include “Secure World” and “Normal World” operating modes of architectures, such as ARM-based TRUSTZONE extensions. It should be understood, however, that OOB MCU **604** described herein is operable with any other security architectures or trust execution environments including, but not limited to Open Multimedia Applications Platform (OMAP)-based architectures, M-Shield-based architectures, x86 system management mode (SMM), etc.

(215) For example, OOB MCU **604** may use TRUSTZONE extensions to enable execution isolation for ensuring code and data integrity and confidentiality and the isolation of operations (and storage) from access by the “Normal World” via a firmware-based processor. In other words, OEM sandbox custom firmware **1902** is read from system firmware (or other source) and placed into protected memory and uses TRUSTZONE extensions and security primitives to provide secure execution isolation within a firmware-based sandbox that can be implemented within ARM-based architectures.

(216) Although one OEM-specific sandbox **1902** is specifically shown in system **1900**, in alternative embodiments any number of such sandboxes may be provided in OOB MCU **604**, each sandbox its own execution area separate and distinct from all other sandboxes.

(217) FIG. **20** is a diagram illustrating an example of method **2000** for handling opaque OOB packets in a sandbox. In some embodiments, method **200** may be performed, at least in part, by OOB MCU **604**.

(218) Particularly, method **2000** begins at **2001**. At **2002**, OOB MCU **604** starts processing OOB packets, for example, received by network device **614** during a G3-sniff operation. At **2003**, OOB MCU **604** determines whether an OOB packet to be processed is opaque.

(219) At **2003**, if the OOB packet is not opaque (e.g., baseline), at **2004** OOB MCU **604** may handle (e.g., decode, route, and/or execute) the OOB packet using OOB baseline firmware **1901**. Conversely, if at **2003** the OOB packet is opaque, at **2005** OOB MCU **604** may handle the OOB packet using OEM sandbox custom firmware **1902**, in a segregated execution environment.

(220) In various embodiments, systems and methods described herein may provide a power domain and/or rail architectures that enable power management in SoC **601**. For example, FIG. **21** is a diagram illustrating an example of power domain architecture **2100** for a heterogeneous computing platform configured to perform OOB packet sniffing operations.

(221) As shown, power domain architecture **2100** may include first power domain **2103** comprising a first set of one or more SoC devices, such that first power domain **2103** is configured to receive a first voltage supply **2102** during an OOB packet sniffing operation while a second set of one or more devices integrated into SoC **601** remains in a low-power state. Power domain architecture **2100** may also include second power domain **2105** comprising the second set of one or more SoC devices. Second power domain **2105** may be configured to receive second voltage supply **2104**.

(222) Each of power domains **2103** and **2105** has its own voltage supply rail **2102** and **2104** provided by voltage regulator(s) **2101**, respectively. Particularly, within SoC **601**, first voltage supply rail **2102** powers OOB MCU **604**, OOB boot ROM **605**, integrated SRAM **606**, network traffic routing controller **610**, and high-bandwidth bus controller **613**. Outside of SoC **601**, first voltage supply rail **2102** also powers network device **614**.

(223) Meanwhile, within SoC **601**, second voltage supply rail **2105** powers host processor(s) **602**, low-bandwidth bus controller **611**, and memory controller **607**. Outside of SoC **601**, second voltage supply rail **2105** also powers soldered down memory **608**, DIMMs **609**, flash memory **612**, and EC **109**.

(224) In various implementations, second voltage supply rail **2105** may be capable of providing more power than first voltage supply rail **2103**. For example, first voltage supply rail **2103** may be configured to supply enough power for network device **614** and OOB MCU **604** to perform G3 Sniff operations (and other operations associated with sniffing), whereas second voltage supply rail **2105** may be configured to supply runtime power sufficient for host processor(s) **602** to execute BIOS/UEFI **107**, load host OS **300**, etc.

(225) In some cases, second power domain **2105** may be configured to receive second voltage supply **2104** under control of network device **614**. For example, network device **614** may include a GPIO pin or the like that controls the status of a power control FET or switch to couple or decouple second voltage rail **2104** from second power domain **2105**. Additionally, or alternatively, first voltage supply rail **2102** may be turned on without second voltage supply rail **2104**; but second voltage supply rail **2104** cannot be turned on without first voltage supply rail **2102** also being on.

(226) In some embodiments, systems and methods described herein may provide progressive power level changes (at system level) so simple OOB packets or commands that can be processed and/or executed efficiently, without waking up unnecessary devices. For example, OOB packets or commands requiring only OOB MCU **604** may not need to wake EC **109**. Similarly, OOB packets or commands requiring only EC **109** may not need to wake up BIOS/UEFI **107** (e.g., EC telemetry for hardware diagnostics). OOB packets or commands requiring BIOS/UEFI **107** may cause SoC **601** to enter its highest power state and bring up all power rails necessary for the host processor(s) **602** to run UEFI.

(227) In that regard, FIG. **22** is a diagram illustrating an example of system **2200** for progressive power control in a heterogeneous computing platform. In this case, SoC **601** comprises: network device **614** disposed in first power domain **2203** configured to receive power from first voltage supply rail **2202** from voltage regulator **2201**, OOB MCU **604** disposed in second power domain **2205** configured to receive power from second voltage supply rail **2204** from voltage regulator **2201**, EC **109** disposed in third power domain **2207** configured to receive power from first voltage supply rail **2206** from voltage regulator **2201**, and host processor(s) **602** disposed in fourth power domain **2209** configured to receive power from fourth voltage supply rail **2208** from voltage regulator **2201**.

(228) In operation, network device **614** examines an OOB packet and powers second power domain **2205** before it routes the OOB packet to OOB MCU **604**. If OOB MCU **604** determines that the OOB packet is opaque, it powers third power domain **2207** and sends the opaque OOB packet to EC **109**. If EC **109** decodes the opaque OOB packet with an OEM's map, key, or decode ring and interprets an OOB management command to be performed by BIOS/UEFI **107** (e.g., change BIOS settings), it powers fourth power domain **2209**.

(229) In some implementations, network device **614** may include a GPIO pin or the like that controls the status of a power control FET or switch to couple or decouple second voltage rail **2104** from second power domain **2105**. OOB MCU **604** may include a GPIO pin or the like that controls the status of another power control FET or switch to couple or decouple third voltage rail **2106** from third power domain **2107**. Furthermore, EC **109** may include a GPIO pin or the like that controls the status of yet another power control FET or switch to couple or decouple fourth voltage rail **2108** from fourth power domain **2108**.

(230) FIG. **23** is a diagram illustrating an example of method **2300** for progressive power control in a heterogenous computing platform. In various embodiments, devices within SoC **601** may be powered up sequentially, depending upon what needs to be done next. For example, network device **614** may be configured bring up EC **109** or OOB MCU **604**, but not host processor(s) **602**.

(231) Particularly, method **2300** starts at **2301**. At **2302**, method **2300** may power up first power domain **2202** comprising network device **614** (e.g., upon expiration of RTC timer **901**). (while second, third, and fourth power domains **2205**, **2207**, and **2209** remain in G3).

(232) At **2303**, method **2300** may determine whether an OOB packet or command retrieved by network device **614** as the result of a G3 Sniff operation is to be processed by OOB MCU **604**. If so, at **2304** network device **614** may power up second power domain **2205** and send the OOB packet or command to OOB MCU **604** (while third and fourth power domains **2207** and **2209** remain in G3).

(233) At **2305**, method **2300** may determine whether the OOB packet or command processed by OOB MCU **604** is to be further processed by EC **109**. If so, at **2306** OOB MCU **604** may power up third power domain **2207** and send the OOB packet or command to EC **109** (while fourth power domain **2209** remains in G3).

(234) At **2307**, method **2300** may determine whether the OOB packet or command processed by EC **109** is to be further processed by host processor(s) **602**. If so, at **2308** OOB MCU **604** may power up fourth power domain **2209** and send the OOB packet or command to host processor(s) **602**.

(235) At **2309**, method **2300** completes the OOB management operations triggered by the OOB packet or command, powers all power domains down (e.g., to G3), and runs RTC timer **901**. At **2310**, upon expiration of RTC timer **901**, control returns to **2302**.

(236) In some cases, operations **2305** and **2306** may be switched with operations **2303** and **2304**, depending upon which device is supposed to decode the OOB packet or execute an OOB command. For example, based upon destination endpoint information (e.g., identifying OOB MCU **604** or EC **109**) contained in an opaque OOB packet header, the EC domain (third power domain **2207**) may be powered up before the OOB MCU domain (second power domain **2207**).

(237) In some embodiments, systems and methods described herein may provide an interconnect architecture suitable for OOB manageability in SoC **601**. Using these systems and methods, OOB MCU **604** may access other integrated devices through SoC **601**'s internal interconnect fabric (e.g., interconnect **203** in FIG. **2**.) while host processor(s) **602** remains in a low-power state (e.g., G3). In some cases, the SoC **601**'s interconnect architecture may provide OOB MCU **604** with access to all PCIe devices, storage, video, etc.

(238) FIG. **24** is a diagram illustrating an example of interconnect architecture **2400** suitable for OOB manageability in a heterogenous computing platform. In this case, OOB MCU **604** is coupled to a plurality of devices **2402-2411** integrated into SoC **601** via internal interconnect **2401** (e.g., **203**).

(239) Specifically, OOB MCU **604** is coupled via interconnect **2401** to integrated GPU (iGPU) **2402** (e.g., **204**), USB controller **2403** (e.g., **214**), storage device or Serial ATA Attachment (SATA)/Nonvolatile Memory Express (NVMe) controller **2404**, sensor hub **2405** (e.g., **207**), audio controller **2406** (e.g., **206**), VPU **2407** (e.g., **211**), I.sup.2C/I.sup.3C controller **2408**, GPIO controller **2409**, network controller **2410** (e.g., **213**), and TPM **2411** (e.g., **212**).

(240) It should be noted that devices **2402-2411** are shown by way of example only. In some cases, one or more of devices **2402-2411** may be absent. In other cases, one or more devices other than devices **2402-2411** may be present.

(241) In some implementations, OOB MCU **604** may be coupled to one or more power control FETs (e.g., **705**), VCI pin wake logic (e.g., **903**), and/or voltage regulator (e.g., **2201**) to control the availability of power to one or more of devices **2402-2411**, individually, as well as that of interconnect **2401**. As such, OOB MCU **604** may have the ability to wake up each device from their low power states, and to set them back into their low-power states, without waking up host processor(s) **602** from its own low power state. Furthermore, once a selected device has been woken up from its low-power state, OOB MCU **604** may send and receive packets, messages, and/or commands to that device through interconnect **2401**.

(242) FIG. **25** is a diagram illustrating an example of method **2500** for using interconnect architecture **2400**, by OOB MCU **604**, to perform manageability operations while host processor(s) **602** remains in a low-power state. Particularly, method **2500** begins at **2501**.

(243) At **2502**, OOB MCU **604** receives an OOB command, for example, included in an OOB packet as a result of a G3 Sniff operation. At **2503**, OOB MCU **604** identifies a selected one of devices **2401-11** to process or execute at least a portion of the OOB command.

(244) For example, OOB MCU **604** may identify the selected device based, at least in part, upon a look up table that associates: (i) commands or types of commands; with (ii) corresponding ones of the plurality of devices or types of devices. Additionally, or alternatively, OOB MCU **604** may identify the selected device based, at least in part, upon a field within a header of the OOB packet within which the command is received (while the host processor is in the low-power state) that includes an identification of the selected device.

(245) In some cases, the selected device may be iGPU **2402**, and the command may include a request to retrieve contents of a video buffer (e.g., while host processor(s) **602** is in its low-power state). Additionally, or alternatively, the selected device may include sensor hub **2405**, and the command may include a request to identify a proximity or presence state of a user of IHS **100** using a proximity sensor (e.g., while host processor(s) **602** is in its low-power state).

(246) Additionally, or alternatively, the selected device may include audio controller **2406**, and the command may include a request to retrieve an audio input using a microphone (e.g., while host processor(s) **602** is in its low-power state). Additionally, or alternatively, the selected device may include a camera controller, and the command may include a request to retrieve a video input using a camera (e.g., while host processor(s) **602** is in its low-power state).

(247) Additionally, or alternatively, the selected device may include storage device controller **2404** (e.g., a hard drive), and the command may include a request to retrieve contents of the storage device (e.g., while host processor(s) **602** is in its low-power state). Additionally, or alternatively, the selected device may include VPU **2407** and the command may include a request to perform an AI/ML operation (e.g., while host processor(s) **602** is in its low-power state).

(248) At **2504**, OOB MCU **604** may power up the selected device out of a low-power state via one or more power control FETs, VCI pin wake logic, and/or a voltage regulator while host processor(s) **602** is in the low-power state. At **2505**, OOB MCU **604** may send the command or an indication of the command to the selected device through internal interconnect **2401**.

(249) At **2506**, OOB MCU **604** determines whether a response to the command or instruction has been received from the selected device via interconnect **2401** indicating that the command or instruction has been processed. If so, at **1207** OOB MCU **604** may power down the selected device to its low-power state via the one or more power control FETs, VCI pin wake logic, and/or voltage regulator.

(250) In various embodiments, systems and methods described herein may provide OOB MCU **604** and/or EC **109** with access to a local storage device (e.g., NVMe, SSD, etc.), regardless of which power state IHS **100** and/or SoC **601** are in (e.g., S0, G3, S5, etc.). For example, OOB MCU **604**

may have access to a local storage device using a main, direct channel (e.g., high-bandwidth bus controller **613**), or a sideband, indirect channel through EC **109** (e.g., via low-bandwidth bus controller **611** and/or interconnect fabric **203**).

(251) In some cases, OOB MCU **604** may also have access the local storage device by bringing up BIOS/UEFI **107** silently, without booting up host OS **300** (ordinarily, once host OS **300** is initialized, OOB MCU **604** loses access to the local storage device in favor of host OS **300**).

(252) As such, these systems and methods may enable multiple OOB manageability use cases, such as, for example, upgrading host OS **300**'s components, image wipe, reimaging host OS **300**, access to crash dumps for diagnostics, etc. In some cases, access to the local storage device may be enabled by a baseline firmware of OOB MCU **604**, and in other cases by a sandbox firmware of OOB MCU **614** (e.g., OEM-specific features, including the processing of opaque OOB packets and commands).

(253) FIG. **26** is a diagram illustrating an example of system **2400** for local storage access by OOB processors in a heterogenous computing platform. As shown, EC **109** is coupled to storage device **2601** via sideband interface **2602**. OOB MCU **604** has direct access to storage device **2601** through main interface **2603** via high-bandwidth bus controller **613** (or interconnect **203**), and to EC **109** via low-bandwidth controller **611**. As such, OOB MCU **604** also has access to storage device **2601**, through EC **109**, via sideband interface **1602**.

(254) Storage device **2601** may include various partitions, sections, or portions, such as, for example: NVMe boot partition **2604**, NVMe protected partitions **2605**, host OS **300**'s boot files **2606**, host OS **300**'s kernel files **2607**, host OS **300**'s application files **2608**, user files **2609**, and security files **2610**. It should be noted that other components, such as network device **614**, for example, are not explicitly shown in system **2400** but may nonetheless still be used.

(255) In operation, network device **614** may receive an OOB command that requires access to storage device **2601**, and it may forward that OOB command to OOB MCU **604**. If host processor(s) **602** are in an S0 state such that host OS **300** is running, OOB MCU **604** may have arbitrated, direct access (through network traffic routing controller **610** and high-bandwidth bus controller **613** via main channel **2603**) to local storage device **2601** in order to execute the OOB command. Additionally, or alternatively, if host processor(s) **602** is in S0 but only BIOS/UEFI **107** is running, or if host processor(s) **602** is in G3 or S5, OOB MCU **604** may have both: direct access to local storage device **2601**, and indirect access (through low-bandwidth bus controller **611** and EC **109** via sideband channel **2602**) to local storage device **2601** in order to execute the OOB command.

(256) FIG. **27** is a diagram illustrating an example of method **2700** for local storage access by OOB processors in heterogenous computing platforms. In various embodiments, method **2500** may be performed by OOB MCU **604** and/or EC **109** in cooperation with network device **614**.

(257) Particularly, method **2700** starts at **2701**. At **2702**, OOB MCU **604** receives an OOB packet from ITDM terminal **618** via network device **614**. At **2703**, network device **614** identifies an OOB command **2401** or type of OOB command within the OOB packet.

(258) At **2704**, OOB MCU **604** identifies (e.g., using a lookup table of OOB commands) whether the OOB command requires read/write access to local storage device **2601**, such as a command to read or write from any of: NVMe boot partition **2604**, NVMe protected partitions **2605**, host OS **300**'s boot files **2606**, host OS **300**'s kernel files **2607**, host OS **300**'s application files **2608**, user files **2609**, or security files **2610**. In response to the identification, OOB MCU **604** may also wake local storage device **2601** from a low-power or off mode (e.g., using a GPIO signal coupled to a VR, or the like), or it may request that EC **109** wake local storage device **2601**.

(259) If so, assuming host processor(s) **602** is in a low-power mode, at **2705** OOB MCU **604** may forward a read/write indication to EC **109** in order to execute the OOB command at **2706** via sideband channel **2602** while host processor(s) **602** remains in the low-power state. If not, control passes from **2704** to **2706**.

(260) In other situations, however, such as when BIOS/UEFI 107 is executed by host processor(s) 602 without host OS 300 running, OOB MCU 604 may access local storage device 2601 directly through main channel 2603, without relying upon EC 109 and/or sideband channel 2602.

Moreover, if host OS 300 is executing, directed access by OOB MCU 604 to local storage device 2601 may be arbitrated by network traffic routing controller 610.

(261) In some cases, once the instructions have been performed and the OOB command executed, OOB MCU 604 may return local storage device 2601 to an off or low-power state (e.g., using a GPIO signal coupled to a VR, or the like).

(262) To implement various operations described herein, computer program code (i.e., program instructions for carrying out these operations) may be written in any combination of one or more programming languages, including an object-oriented programming language such as Java, Smalltalk, Python, C++, or the like, conventional procedural programming languages, such as the “C” programming language or similar programming languages, or any of machine learning software. These program instructions may also be stored in a computer readable storage medium that can direct a computer system, other programmable data processing apparatus, controller, or other device to operate in a particular manner, such that the instructions stored in the computer readable medium produce an article of manufacture including instructions which implement the operations specified in the block diagram block or blocks.

(263) Program instructions may also be loaded onto a computer, other programmable data processing apparatus, controller, or other device to cause a series of operations to be performed on the computer, or other programmable apparatus or devices, to produce a computer implemented process such that the instructions upon execution provide processes for implementing the operations specified in the block diagram block or blocks.

(264) Modules implemented in software for execution by various types of processors may, for instance, include one or more physical or logical blocks of computer instructions, which may, for instance, be organized as an object or procedure. Nevertheless, the executables of an identified module need not be physically located together but may include disparate instructions stored in different locations which, when joined logically together, include the module and achieve the stated purpose for the module. Indeed, a module of executable code may be a single instruction, or many instructions, and may even be distributed over several different code segments, among different programs, and across several memory devices.

(265) Similarly, operational data may be identified and illustrated herein within modules and may be embodied in any suitable form and organized within any suitable type of data structure.

Operational data may be collected as a single data set or may be distributed over different locations including over different storage devices.

(266) Reference is made herein to “configuring” a device or a device “configured to” perform some operation(s). It should be understood that this may include selecting predefined logic blocks and logically associating them. It may also include programming computer software-based logic of a retrofit control device, wiring discrete hardware components, or a combination of thereof. Such configured devices are physically designed to perform the specified operation(s).

(267) It should be understood that various operations described herein may be implemented in software executed by processing circuitry, hardware, or a combination thereof. The order in which each operation of a given method is performed may be changed, and various operations may be added, reordered, combined, omitted, modified, etc. It is intended that the invention(s) described herein embrace all such modifications and changes and, accordingly, the above description should be regarded in an illustrative rather than a restrictive sense.

(268) Unless stated otherwise, terms such as “first” and “second” are used to arbitrarily distinguish between the elements such terms describe. Thus, these terms are not necessarily intended to indicate temporal or other prioritization of such elements. The terms “coupled” or “operably coupled” are defined as connected, although not necessarily directly, and not necessarily

mechanically. The terms “a” and “an” are defined as one or more unless stated otherwise. The terms “comprise” (and any form of comprise, such as “comprises” and “comprising”), “have” (and any form of have, such as “has” and “having”), “include” (and any form of include, such as “includes” and “including”) and “contain” (and any form of contain, such as “contains” and “containing”) are open-ended linking verbs.

(269) As a result, a system, device, or apparatus that “comprises,” “has,” “includes” or “contains” one or more elements possesses those one or more elements but is not limited to possessing only those one or more elements. Similarly, a method or process that “comprises,” “has,” “includes” or “contains” one or more operations possesses those one or more operations but is not limited to possessing only those one or more operations.

(270) Although the invention(s) is/are described herein with reference to specific embodiments, various modifications and changes can be made without departing from the scope of the present invention(s), as set forth in the claims below. Accordingly, the specification and figures are to be regarded in an illustrative rather than a restrictive sense, and all such modifications are intended to be included within the scope of the present invention(s). Any benefits, advantages, or solutions to problems that are described herein with regard to specific embodiments are not intended to be construed as a critical, required, or essential feature or element of any or all the claims.

Claims

1. A heterogeneous computing platform, comprising: an Out-of-Band (OOB) Microcontroller Unit (MCU) configured to periodically wake from an OOB MCU low-power state in response to a trigger signal originating within the OOB MCU; and a memory coupled to the OOB MCU, the memory having program instructions stored thereon that, upon execution by the OOB MCU, cause the OOB MCU to perform operations comprising: upon wake from the OOB MCU low-power state, receive an OOB command, while every host processor of the heterogeneous computing platform remains in a host processor low-power state; in response to an indication that the OOB command requires a read or write operation, wake up a local storage device external to the heterogeneous computing platform from a local storage device low-power state through an Embedded Controller (EC) coupled to the OOB MCU; and access the local storage device through the EC.
2. The heterogeneous computing platform of claim 1, further comprising: a System-On-Chip (SoC), a Field-Programmable Gate Array (FPGA), or an Application-Specific Integrated Circuit (ASIC).
3. The heterogeneous computing platform of claim 1, further comprising a Reduced Instruction Set Computer (RISC) processor and a plurality of devices coupled to an interconnect.
4. The heterogeneous computing platform of claim 3, wherein the heterogeneous computing platform comprises at least one of: a Graphical Processing Unit (GPU), an audio Digital Signal Processor (aDSP), a sensor hub, a Neural Processing Unit (NPU), a Tensor Processing Unit (TPU), a Neural Network Processor (NNP), an Intelligence Processing Unit (IPU), an Image Signal Processor (ISP), or a Video Processing Unit (VPU).
5. The heterogeneous computing platform of claim 3, wherein the interconnect comprises at least one of: an Advanced Microcontroller Bus Architecture (AMBA) bus, a QuickPath Interconnect (QPI) bus, or a HyperTransport (HT) bus.
6. The heterogeneous computing platform of claim 1, wherein the local storage device comprises at least one of a Non-Volatile Memory Express (NVMe) device or a Solid-State Drive (SSD).
7. The heterogeneous computing platform of claim 1, wherein the EC is coupled to the local storage device via a sideband interface.
8. The heterogeneous computing platform of claim 7, wherein the sideband interface comprises an Inter-Integrated Circuit (I.sup.2C) bus or an Improved Inter-Integrated Circuit (I.sup.3C) bus.

9. The heterogeneous computing platform of claim 1, wherein the read or write operation involves at least one of: a Non-Volatile Memory Express (NVMe) boot partition, an NVMe protected additional partition, an Operating System (OS) boot file, an OS kernel file, an OS application file, a user file, or a security file.
 10. The heterogeneous computing platform of claim 1, wherein the OOB MCU is further configured to access the local storage device through a bus controller comprising a Peripheral Component Interconnect Express (PCIe) controller integrated into the heterogeneous computing platform.
 11. An Information Handling System (IHS), comprising: a storage device; and an Out-of-Band (OOB) Microcontroller Unit (MCU) integrated into a heterogeneous computing platform and coupled to a local storage device through a bus controller comprising a PCIe controller, wherein the OOB MCU is configured to periodically wake from an OOB MCU low-power state in response to a trigger signal originating within the OOB MCU, and wherein, upon wake from the OOB MCU low-power state, the OOB is configured to receive an OOB command causing the OOB MCU to wake the local storage device from a storage device low-power state.
 12. The IHS of claim 11, wherein the local storage device comprises a Non-Volatile Memory Express (NVMe) device or a Solid-State Drive (SSD).
 13. The IHS of claim 11, wherein the OOB MCU is further configured to access the local storage device through an Embedded Controller (EC).
 14. The IHS of claim 13, wherein to access the local storage device through the EC, the OOB MCU is coupled to the EC through a bus controller comprising an Inter-Integrated Circuit (I.sup.2C) bus or an Improved Inter-Integrated Circuit (I.sup.3C) bus.
 15. A method, comprising: receiving, at a heterogeneous computing platform of an Information Handling System (IHS), an OOB packet comprising a read or write command with respect to a local storage device coupled to the heterogeneous computing platform; periodically waking an Embedded Controller (EC) from an EC low-power state in response to a trigger signal originating within the EC, wherein the EC is coupled to the heterogeneous computing platform, and processing, by the EC upon wake from the EC low-power state, the read or write command while remains in a low-power every host processor of the heterogeneous computing platform state.
 16. The method of claim 15, wherein the OOB packet is retrieved by a network device coupled to the heterogeneous computing platform in response to an OOB packet sniffing operation while the host processor is in the low-power state.
 17. The method of claim 15, wherein the read or write command comprises an instruction to reimage a host Operating System (OS) executable by the host processor.
-