

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12405801
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Atluri; Aditya Avinash et al.

Scalarization of instructions for SIMT architectures

Abstract

Apparatuses, systems, and techniques to adapt instructions in a SIMT architecture for execution on serial execution units. In at least one embodiment, a set of one or more threads is selected from a group of active threads associated with an instruction and the instruction is executed for the set of one or more threads on a serial execution unit.

Inventors: Atluri; Aditya Avinash (Redmond, WA), Choquette; Jack (Palo Alto, CA), Edwards; Carter (Campbell, CA), Giroux; Olivier (Santa Clara, CA), Kaushik; Praveen Kumar (Bengaluru, IN), Krashinsky; Ronny (Portola Valley, CA), Kulkarni; Rishkul (Austin, TX), Kyriakopoulos; Konstantinos (Weinsberg, DE)

Applicant: NVIDIA Corporation (Santa Clara, CA)

Family ID: 90574113

Assignee: NVIDIA Corporation (Santa Clara, CA)

Appl. No.: 18/105679

Filed: February 03, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240118899 A1	Apr. 11, 2024

Foreign Application Priority Data

GR	20220100820	Oct. 06, 2022
----	-------------	---------------

Publication Classification

Int. Cl.: G06F9/38 (20180101)

U.S. Cl.:

CPC **G06F9/3851** (20130101);

Field of Classification Search

CPC: G06F (9/3851)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6076155	12/1999	Blomgren	712/225	G06F 9/30123
2003/0145190	12/2002	Faraboschi	712/E9.05	G06F 9/3842
2007/0067607	12/2006	Jiao	712/215	G06F 9/3851
2013/0042090	12/2012	Krashinsky	712/214	G06F 9/38885
2016/0371093	12/2015	Chang	N/A	G06F 9/3851

Primary Examiner: Abad; Farley

Attorney, Agent or Firm: Lowenstein Sandler LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS

(1) This application claims the benefit of priority to Greek Patent Application No. 20220100820, entitled, “SCALARIZATION OF INSTRUCTIONS FOR SIMT ARCHITECTURES,” filed on Oct. 6, 2022, which is incorporated by reference herein in its entirety.

TECHNICAL FIELD

(2) Embodiments of the disclosure generally relate to parallel processing architectures, and more specifically, to improved techniques for executing instructions in a single instruction multiple thread (SIMT) processing architecture.

BACKGROUND

(3) Many computer applications can be accelerated through the use of parallel processing techniques, e.g., where the same instructions can be executed on multiple data elements in parallel. In image and media processing applications, for example, the processing of large sets of pixels, image blocks, and/or vertices can be mapped to different computing threads or processing lanes that can be executed in parallel. For instance, in a single instruction multiple thread (SIMT) processing architecture, a common instruction (or instruction stream) can be executed using a group of processing threads in parallel.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The disclosure will be understood more fully from the detailed description given below and from the accompanying drawings of various embodiments of the disclosure. The drawings, however, should not be taken to limit the disclosure to the specific embodiments, but are for

explanation and understanding only.

- (2) FIG. 1 illustrates an example computing environment, according to at least one embodiment;
- (3) FIG. 2 illustrates a flow diagram of an example method for executing an instruction on a serial execution unit using a collect peeling loop, according to at least one embodiment;
- (4) FIG. 3 illustrates a flow diagram of an example method for executing an instruction on a serial execution unit using a select peeling loop, according to at least one embodiment;
- (5) FIG. 4 illustrates a computer system, according to at least one embodiment;
- (6) FIG. 5 illustrates a parallel processing unit (“PPU”), according to at least one embodiment;
- (7) FIG. 6 illustrates a general processing cluster (“GPC”), according to at least one embodiment;
- (8) FIG. 7 illustrates a memory partition unit of a parallel processing unit (“PPU”), according to at least one embodiment;
- (9) FIG. 8 illustrates a streaming multi-processor, according to at least one embodiment;
- (10) FIG. 9 illustrates a computer system, according to at least one embodiment;
- (11) FIG. 10A illustrates a parallel processor, according to at least one embodiment;
- (12) FIG. 10B illustrates a partition unit, according to at least one embodiment;
- (13) FIG. 10C illustrates a processing cluster, according to at least one embodiment; and
- (14) FIG. 10D illustrates a graphics multiprocessor, according to at least one embodiment.

DETAILED DESCRIPTION

(15) The systems and methods described herein may be used by, without limitation, non-autonomous vehicles, semi-autonomous vehicles (e.g., in one or more adaptive driver assistance systems (ADAS)), piloted and un-piloted robots or robotic platforms, warehouse vehicles, off-road vehicles, vehicles coupled to one or more trailers, flying vessels, boats, shuttles, emergency response vehicles, motorcycles, electric or motorized bicycles, aircraft, construction vehicles, underwater craft, drones, and/or other vehicle types. Further, the systems and methods described herein may be used for a variety of purposes, by way of example and without limitation, for machine control, machine locomotion, machine driving, synthetic data generation, model training, perception, augmented reality, virtual reality, mixed reality, robotics, security and surveillance, simulation and digital twinning, autonomous or semi-autonomous machine applications, deep learning, environment simulation, object or actor simulation and/or digital twinning, data center processing, conversational AI, light transport simulation (e.g., ray-tracing, path tracing, etc.), collaborative content creation for 3D assets, cloud computing and/or any other suitable applications.

(16) Disclosed embodiments may be comprised in a variety of different systems such as automotive systems (e.g., a control system for an autonomous or semi-autonomous machine, a perception system for an autonomous or semi-autonomous machine), systems implemented using a robot, aerial systems, medial systems, boating systems, smart area monitoring systems, systems for performing deep learning operations, systems for performing simulation operations, systems for performing digital twin operations, systems implemented using an edge device, systems incorporating one or more virtual machines (VMs), systems for performing synthetic data generation operations, systems implemented at least partially in a data center, systems for performing conversational AI operations, systems for hosting real-time streaming applications, systems for presenting one or more of virtual reality content, augmented reality content, or mixed reality content, systems for performing light transport simulation, systems for performing collaborative content creation for 3D assets, systems implemented at least partially using cloud computing resources, and/or other types of systems.

(17) At a high level, in a single instruction multiple thread (SIMT) processing architecture, a common instruction (or instruction stream) can be executed using a group of processing threads in parallel. At a hardware level, the parallel execution of multiple threads is performed using a parallel execution unit, such as an SIMT execution unit (e.g., similar to a traditional vector execution unit). The parallel execution unit, for example, may be able to concurrently perform a variety of different

computational operations (e.g., integer and floating-point arithmetic operations, comparison operations, Boolean operations, etc.). Many of these operations may be deterministic in nature (e.g., such as arithmetic and comparison operations) such that performing the operation on a particular set of values is expected to always produce the same result.

(18) In practice, it is often the case that multiple threads concurrently execute the same deterministic instruction on uniform sets of input operands to generate the same result. For example, in graphics processing or finite element analysis applications, computations may be performed on data (e.g., an image or structural model) that exhibit some degree of uniformity (e.g., spatial uniformity in a portion of the image, or temporal uniformity across a series of images). As another example, a counter of a programming loop executed by multiple threads may be incremented following each loop iteration. In such cases, the redundant execution of the instruction using each thread may unnecessarily increase power consumption and resource utilization.

(19) Implementing and/or performing certain compute operations using a parallel execution unit (e.g., a SIMT execution unit) can also be expensive from a silicon use and timing perspective. For example, because compute logic is replicated for each thread lane in a parallel execution unit, implementing complex operations (e.g., tensor memory access (TMA) operations) may require a significant amount of silicon real estate (e.g., significant number of transistors). Furthermore, certain operations may take a substantial amount of time to execute (e.g., several hundreds or thousands of clock cycles) but may only be executed by a few threads at a time, thereby reducing the overall utilization rate of the parallel execution unit.

(20) Embodiments of the present disclosure address the above-mentioned limitations and/or other limitations of existing architectures by adapting a parallel execution model, such as the SIMT execution model, to perform certain operations using a serial execution unit (e.g., similar to a traditional scalar execution unit), capable of executing a single instruction on a single set of input operands. For instance, where multiple threads execute the same deterministic instruction on a uniform set of input operand values, the instruction may be performed using the serial execution unit, with the result being shared with all threads. In this way, greater power efficiency may be achieved—e.g., because the serial execution unit may perform the computation once, and the threads of the SIMT architecture sharing the same instruction may preserve resources by not being required to perform the computation. The serial execution unit, likewise, may implement and perform operations that are too expensive to be provided for, or executed on, a parallel execution unit. This may not only reduce the size and complexity of a parallel execution unit, but also free up the parallel execution unit to perform other operations.

(21) A “scalarization” process may be employed in embodiments to perform this adaptation, where instructions destined for execution using a parallel execution unit (e.g., an SIMT execution unit) are recast for execution by a serial execution unit. The scalarization process may use peeling loops to partition and unwind the execution of threads in a group to be performed (in serial, as necessary) on a serial execution unit. For example, in scalarizing instructions that can execute on behalf of multiple threads, the peeling loop may operate to partition a group of threads into sub-groups that exhibit dynamic uniformity for (e.g., all) source operand values, with the instruction for individual sub-groups being collectively executed once using a serial execution unit. As another example, when scalarizing an instruction that is more suitable for serial execution (e.g., instructions that are not supported by, or would be too expensive to implement or execute using, a parallel execution unit), the peeling loop may operate to unwind and execute the instruction on a serial execution unit one thread at a time.

(22) The peeling loops may be further optimized, e.g., by a compiler, to simplify (e.g., reduce the number of iterations necessary, or eliminate needless operations within the loop) or flatten the peeling loop entirely in some instances. The compiler, for instance, may perform uniformity analysis to determine whether source operands share the same value among executing threads. The compiler, likewise, may be able to determine the number of threads that are expected to execute an

instruction and simplify the peeling loop accordingly (e.g., eliminating the peeling loop where exactly one thread is expected to execute the instruction).

(23) FIG. 1 is a block diagram illustrating a computing system **100**, according to at least one embodiment. In some embodiments, computing system **100** may be a heterogenous computing system that includes one or more types of computational units, including for example, one or more central processing units (CPUs), one or more graphics processing units (GPUs), one or more digital signal processors (DSPs), one or more data processing units (DPUs), one or more field-programmable gate arrays (FPGAs), and/or one or more application specific integrated circuits (ASICs). As illustrated, for instance, computing system **100** may include a general-purpose processor **110** (e.g., a multicore CPU) and a parallel processor **120** (e.g., a general-purpose GPU (GPGPU)).

(24) General purpose processor **110** may be designed for fast serial processing of program instructions, whereas parallel processor(s) **120** may be designed for highly parallel processing of program instructions (e.g., computational instructions). Parallel processor(s) **120** may operate as a coprocessor to the general-purpose processor **110**, where portions of a computer application (e.g., data-parallel, compute intensive portions of an application) are off-loaded to the parallel processor **120** for execution.

(25) As an illustrative example, computing system **100** may be used to execute a computer application **140**. Computer application **140** may include a collection of program instructions that may include a mix of sequential instruction portions **141**, which may be executed as a series of one or more threads **145** on general purpose processor **110**, and parallel instruction portions **142**, which may be executed in parallel as multiple threads **146** on parallel processor **120**. A portion of computer application, for example, may contain programming instructions that are executed many times, but independently on potentially different data, which can be executed as multiple threads **146** on parallel processor **120**. The threads **146** may be organized as one or more thread blocks **147** (e.g., as an array or grid **148** of thread blocks **147**), which may be concurrently executed by parallel processor **120**. In some embodiments, parallel processor **120** may include one or multiprocessors **130**, with one or more thread blocks **147** being distributed to each multiprocessor **130** for execution. Individual threads in a thread block **147** can be executed concurrently by multiprocessors **130**, and multiprocessors **130** can execute multiple thread blocks concurrently.

(26) In some embodiments, multiprocessors **130** may employ a SIMT (Single-Instruction, Multiple-Thread) architecture for concurrent execution of multiple threads. By way of example, multiprocessors **130** may be configured to create, manage, schedule, and execute threads in groups of parallel threads, which may be referred to as a warp. When multiprocessors **130** are given one or more thread blocks to execute, they may partition them into separate thread groups, which may be independently scheduled for execution.

(27) In some embodiments, multiprocessors **130** may be configured to execute one common instruction for a group of threads at a time (e.g., a warp, half-warp, quarter-warp, etc.). Full efficiency, thus, may be realized when all threads in the group agree on their execution path. Individual threads within a group of threads may start together at a same program address (e.g., a common instruction in a sequence of instructions) but may be assigned their own instruction address counter (or program counter) and register state, allowing each thread to branch and execute independently. If individual threads diverge via a conditional control construct (e.g., a conditional branch, conditional function call, or conditional return), the different branch paths (e.g., resulting from the divergence) may be serially executed. When execution of (e.g., all) branch paths complete, the threads may converge back to the same execution path. In some cases, a program instruction may provide a synchronization point where all threads in the group converge (e.g., where some threads may wait until all threads in the thread arrive). Threads in a group of threads that are participating in the current instruction may be referred to as the active threads, whereas threads not on the current instruction may be referred to as inactive (or disabled) threads.

(28) In some embodiments, multiprocessors **130** may include functional execution units that may be configured to perform a variety of operations, including for example, integer and floating-point arithmetic operations (e.g., addition and multiplication), comparison operations, Boolean operations (AND, OR, XOR), bit-shifting operations, random number generation operations, and other computational operations (e.g., planar interpolation and trigonometric, exponential, and logarithmic functions, etc.). In some embodiments, multiprocessors **130** may include one or more parallel execution unit(s) **131** and serial execution unit(s) **132**. Parallel execution units **131** may be configured to execute a single instruction on multiple sets of data (e.g., similar to a traditional vector execution unit). In some embodiments, parallel execution units **131** may be able to execute a common instruction for each thread in a group of threads (e.g., using a distinct set of source operands and resulting in a distinct set of result operands for each thread). Serial execution units **132** may be configured to execute a single instruction on a single set of data (e.g., similar to a traditional scalar execution unit). In some embodiments, serial execution units **132** may be able to execute an instruction for one or more threads (e.g., using a single set of shared source operands and resulting in a single set of shared destination operands for all threads). The operations supported by parallel execution units **131** and serial execution units **132** may vary, with different embodiments including parallel execution units **131** and serial execution units **132** that commonly support certain operations and/or uniquely support other operations. For example, in some embodiments, parallel execution units **131** and serial execution units **132** may commonly support certain arithmetic operations (e.g., integer addition or multiplication operations) and/or other computational operations. As another example, certain memory operations (e.g., TMA operations) may be particularly complex (e.g., requiring significant amount of silicon to implement) and relatively expensive to implement as a parallel execution unit **131**. Accordingly, in some embodiments, multiprocessors **130** may support execution of such operations by serial execution units **132**.

(29) In some embodiments, multiprocessors **130** may include one or more sets of register files, or registers, for use by the functional execution units of multiprocessors **130**. In some embodiments, for example, multiprocessors **130** may include a set of private registers **133** that may provide temporary storage for operands connected to data paths of parallel execution units **131**. In some embodiments, private registers **133** may be partitioned and allocated to individual threads in a group of threads being executed by multiprocessors **130**, with the allocated portion serving as a private register space of each of the individual threads. In some embodiments, private registers **133** may be statically partitioned (e.g., having a fixed size for each individual thread) and dynamically allocated for use by individual threads. In some embodiments, multiprocessors **130** may include a set of shared registers **134** that may provide temporary storage for operands connected to data paths of serial execution units **132**. In some embodiments, shared registers **134** may be accessible by some or all threads in a group of threads being executed by multiprocessors **130**.

(30) In some embodiments, multiprocessor **130** may include one or more additional sets of registers. In some embodiments, for example, multiprocessor **130** may include a set of special registers that may store predefined, platform-specific information, such as thread parameters (e.g., a thread identifier (within a thread block), lane identifier (within a warp), warp identifier, block identifier, etc.), clock counters, and/or performance monitoring information. In some embodiments, multiprocessor **130** may include a set of predicate registers **135** that may be used to store predicates (e.g., a 1-bit Boolean value), which may be used to support instruction predication (e.g., conditional branch predication). In some embodiments, for example, an instruction may accept an optional predicate guard operand, which if determined to be true, may cause the instruction to be executed and if determined to be false, may preclude execution of the instruction. In some embodiments, a predicate mask (or predicate vector) may be used as a predicate guard for an instruction to be executed by multiple threads, e.g., with each element of the mask corresponding to a particular thread lane. In some embodiments, for example, a predicate mask may be used to

identify active threads in a group of threads (e.g., that are participating in the current instruction) and inactive (or disabled) threads in the group of threads (e.g., that are not performing the current instruction).

(31) In some embodiments, multiprocessors **130** may include a local memory **136** for use by the functional execution units of multiprocessors **130**. In some embodiments, local memory **136** may include a private local memory space that may be (statically or dynamically) allocated to and accessed by individual threads in a group of threads being executed by multiprocessors **130**, with the allocated portion serving as a private memory space of the individual thread. In some embodiments, local memory **136** may also include a shared memory space that may be (statically or dynamically) allocated to and accessed by some or all threads in a group of threads being executed by multiprocessors **130**. In some embodiments, multiprocessors **130** may also be able to access a global memory space, e.g., on a device memory **121** of parallel processor **120**, which may be provided to some or all threads in a group of threads being executed by multiprocessors **130**.

(32) In some embodiments, computing system **100** may include software compiler logic **150** that may be used to compile a computer application (e.g., computer application **140**) from program code, which may be stored in memory **112**. Software compiler logic **150**, for example, may be used to compile program source code **151** into binary code **154** that may be executed by general purpose processor **110** and/or parallel processor **120**. Program source code **151** may include a mix of code, some of which may be designed to execute on general purpose processor **110** (“host code”) and some of which may be designed to execute on parallel processor **120** (“device code”). In some embodiments, software compiler logic **150** may operate to separate device code from host code and compile the code separately. Software compiler logic **150**, for instance, may compile device source code, e.g., into one or more function kernels, and then modify the host code to include the necessary runtime functional calls to load and launch each compiled function kernel. Software compiler logic **150** may then compile the modified host code to obtain binary code **154**, which may be executable by computing system **100** on general purpose processor **110** and parallel processor **120**.

(33) In some embodiments, software compiler logic **150** may compile program source code **151**, or a portion thereof (e.g., device code in program source code **151**), in multiple stages, generating one or more sets of intermediate code (e.g., intermediate assembly code **152** and low-level assembly code **153**) before ultimately arriving at binary code **154**. Program source code **151**, for instance, may be written using a high-level programming language (e.g., C, C++, Java, Python, Fortran, DirectCompute, OpenACC, etc.). Software compiler logic **150** may compile program source code **151** written in a high-level programming language into intermediate assembly code **152** (e.g., PTX code, Khronos SPIR code, LLVM IR code, etc.), which may use a particular instruction set architecture (ISA). In some embodiments, intermediate assembly code **152** may use an instruction set suitable for general purpose parallel programming, which may be designed for efficient execution by parallel processors **120**. In some embodiments, software compiler logic **150** may compile device code in program source code **151** into an intermediate assembly code **152** that is designed to be architecture independent, so the same code can be used for different parallel processor architectures. In some embodiments, a computer application (or portion thereof) may be directly written as intermediate assembly code **152**. In some embodiments, software compiler logic **150** may operate to translate (e.g., further compile) intermediate assembly code into low-level assembly code **153** (e.g., Source and Assembly (SASS) code). Low-level assembly code **153** may use another ISA (e.g., distinct from that of intermediate assembly code **152**), which may be a native architecture that uses target-architecture instructions for particular parallel processor architectures. In some embodiments, software compiler logic **150** may use low-level assembly code **153** to generate binary microcode for native execution on a parallel processor **120**.

(34) In some embodiments, software compiler logic **150** may operate to generate (e.g., optimized) program code (e.g., intermediate assembly code **152**, low-level assembly code **153**, or binary code

154), which for example, may improve execution efficiency and resource utilization. Software compiler logic **150**, for example, may seek to optimize program code to expose sufficient parallelism, coalesce memory access, ensure coherent execution within a group of threads, etc., which may improve execution of the program code on parallel processor **120** (and multiprocessors **130**).

(35) In some embodiments, for example, software compiler logic **150** may operate to perform branch predication (or control flow flattening) to ensure convergent execution of multiple threads (e.g., all threads in a group of threads). In some embodiments, for example, software compiler logic **150** may analyze program code and determine instances where branch predication may be used to implement the code more efficiently. Software compiler logic **150**, for instance, may determine when program loops (e.g., for, while, do-while loops) or logic blocks (e.g., if or switch blocks) create branches in the program code that may be cheaper to predicate and execute for all threads (e.g., cheaper than branching the code and serially executing each branch path) and may unroll or flatten these loops or logic blocks. That is, instructions whose execution depends on a conditional control construct (e.g., a conditional branch, conditional function call, or conditional return) are not skipped, but instead are associated with a per-thread condition code, or predicate, that is set to true or false based on the controlling condition. While these instructions may be scheduled for execution by all threads, only those instructions (or threads) with a true predicate value are actually executed. For instructions (or threads) having a false predicate value, addresses may not be evaluated, operands may not be read, and/or results may not be written. Predicating and executing the instruction for all threads may be cheaper.

(36) In some embodiments, software compiler logic **150** may generate code that is optimized for execution by a particular parallel processor **120**, e.g., taking into account the different functional execution units of the parallel processor **120** (and its multiprocessors **130**) and the operations they support. For example, in a SIMT programming and execution model, instructions may be executed by multiple threads concurrently, with each thread reading potentially distinct input data in respective source operands and generating potentially distinct results to respective destination operands. While instructions, at a hardware level, may generally be executed using a parallel execution unit **131**, it may not always be possible or advantageous to do so, and some instructions may need to be, or may preferably be, executed using a serial execution unit **132**. Software compiler logic **150** may detect such instances (e.g., where execution of an instruction using a serial execution unit **132** is necessary and/or preferable) and may operate to adapt program code written for parallel execution to be executed on serial execution units **132**. Software compiler **150** may further operate to optimize the adapted code for more efficient execution (e.g., by eliminating unnecessary operations, performing branch predication, etc.).

(37) In some instances, multiple threads may concurrently execute the same deterministic instruction on uniform sets of input operands to generate the same result. In such cases, the redundant execution of the instruction by each thread (e.g., using a parallel execution unit **131**), unnecessarily increases power consumption and resource utilization. Therefore, in some embodiments, software compiler logic **150** may be able to analyze program code to determine whether an instruction that is to be executed by multiple threads exhibits execution uniformity (e.g., where the threads are to execute the instruction concurrently) and operand uniformity (e.g., where the source operands of the instruction for the threads have the same value). In some embodiments, for example, software compiler logic **150** may perform iterative flow analysis to determine whether an instruction exhibits execution and/or operand uniformity. Based on this analysis, software compiler logic **150** may determine whether to execute the instruction or serial execution units **132** (e.g., for improved execution efficiency).

(38) In some cases, software compiler logic **150** may be able to determine with certainty that a deterministic instruction to be executed by a group of active threads exhibits complete execution uniformity (e.g., for all threads) and operand uniformity (e.g., for all source operands). Software

compiler logic **150**, for example, may be able to determine that the instruction is to be executed by a single thread, based on which software compiler logic **150** may conclude that the instruction exhibits complete operand uniformity (e.g., the single thread is uniform with itself). In such cases, instead of generating an instruction to be executed by a parallel execution unit **131** (“parallel instruction”), software compiler logic **150** may generate an instruction (or several instructions) that affects execution of the instruction by a serial execution unit **132** (“uniform instruction”). Software compiler logic **150**, for example, may generate instructions to copy any source operand values stored in private registers **133** into shared registers **134** (e.g., to obtain shared source operands accessible by a serial execution unit **132**) and a uniform instruction to be executed using the shared source operands. In some embodiments, software compiler logic **150** may generate the parallel instruction in the first instance, and then replace the parallel instruction with the uniform instruction (or uniform instructions).

(39) As an illustrative example, program source code **151** may include an addition operation, in which three integer values are to be added together, to be performed for each thread in a group of threads. Software compiler logic **150** may typically generate the following assembly instruction (e.g., in low-level assembly code **153**), which may operate to perform the addition operation separately for each thread using a parallel execution unit **131**: `iadd r9, r2, r5`;

where `iadd` may be a parallel instruction that adds the values of source operands `r2` and `r5`, which may be private registers (e.g., in the register space of each thread), and places the result in destination operand `r9`, which may be a private register (e.g., in the register space of each thread).

(40) Where software compiler logic **150** is able to determine that the source operand values (e.g., values in private registers `r2` and `r5`) will be the same for each thread, it may generate the following assembly instructions instead, to affect execution of the addition operation once on behalf of all threads using a serial execution unit **132**:

(41) `TABLE-US-00001 mov ur2, r2; mov ur5, r5; uiadd ur9, ur2, ur5`;

where the `mov` instructions may operate to copy source operand values from private registers `r2` and `r5` (e.g., from a register space of a particular thread) to shared registers `ur2` and `ur5` (e.g., to the shared register space accessible by all threads), and `uiadd` may be a uniform instruction that performs the addition operation once using the shared source operands on serial execution unit **132**. The result of the add operation may, optionally, be copied back to a private register for each thread, e.g., as follows: `mov r9, ur9`;

(42) In other cases, software compiler logic **150** may be able to determine that there is a high likelihood (but not certainty) that a deterministic instruction to be executed by a group of active threads exhibits execution uniformity and operand uniformity. In such cases, software compiler logic **150** may generate a peeling loop that may operate to partition and unwind execution of the group of threads into one or more sets of threads having shared source operand values that are executed once (e.g., for all threads in a particular set of threads) in serial fashion (e.g., for each of the one or more set of threads) using serial execution unit **132**.

(43) In some embodiments, software compiler logic **150** may make use of a collect instruction that may operate to collect threads in a group of threads having a common source operand value and, as necessary, copy the common source operand value from a private register **133** to a shared register **134**. In some embodiments, the collect instruction may be executed by parallel execution unit **131** as a collective operation. That is, parallel execution unit **131** in executing the collect instruction may perform a collective operation in which the group of threads executing the instruction work together to choose a common source operand value and collect those threads having that common source operand value.

(44) In some embodiments, the collect instruction may take the following form: `@pg collect pd, urd, rs, ps`

The collect instruction may have two source operands `rs`, `ps`, two destination operands `urd`, `pd`, and an optional predicate guard `@pg`. Source operand `ps` may be a predicate mask identifying a group

of active threads for which an instruction is pending execution (e.g., for which the predicate mask ps is true) from which a set of threads is to be collected. The instruction may operate to identify a set of threads where the value of private register rs is the same value and copy the value to shared register urd. A destination predicate pd may be set to true for the identified threads, with any other active threads being set to false. In some embodiments, the collect instruction may be fully predicated using predicate guard @pg.

(45) In some embodiments, for example, software compiler logic **150** may generate a “collect” peeling loop, where each iteration of the peeling loop may include collect instructions that are recursively called to collect a set of threads having uniform source operand values (e.g., where each source operand of the instruction has the same value across all threads) and place the source operand values in shared registers **134** (as necessary). A uniform instruction may then be executed by serial execution unit **132** using the uniform source operand values, the result of which may, optionally, be copied back to a private register for threads in the set of threads for which the instruction was executed. The predicate mask identifying active threads may be updated (e.g., to exclude those threads for which the instruction was just executed) and the peeling loop may proceed to the next iteration, ultimately terminating when no threads remain pending execution.

(46) As an illustrative example, software compiler logic **150** may generate the following peeling loop for the addition operation described above:

(47) TABLE-US-00002 set p1, pTrue // 1 Ltop: collect p0, ur5, r5, p1 // 2 collect p0, ur2, r2, p0 // 3 uiadd ur9, ur2, ur5 // 4 @p0 mov r9, ur9 // 5 @p1 mov p1, ~p0 // 6 @p1 branch.u.any Ltop // 7

The predicate p1 may be used to store the active threads for which the instruction is pending execution. As illustrated above, predicate p1 is initialized to true (by pTrue), e.g., for all threads in a group of threads (e.g., in a warp), but in other cases predicate p1 may inherit the predication of the instruction being executed, e.g., where the addition operation is to be performed for a subset of a group of threads (e.g., a half warp, a quarter warp, etc.).

(48) The loop may begin with a first collect instruction that may operate to collect a set of threads having the same value in a first source operand, e.g., private register r5, placing the value into shared register ur5 and setting predicate pG to true for those threads that were just collected (e.g., with all other active threads set to false). In some embodiments, the selection of an anchor source operand value (e.g., based on which the collection is performed) may be identified at random from amongst the threads being collected, while in other embodiments, the source operand value of a first thread (e.g., having a lowest thread ID) may be used as the anchor value for the collect operation.

(49) The collect operation may be repeated, recursively, for the remaining source operands. A second collect operation, for instance, may operate to collect a subset of threads from those previously collected (e.g., from the threads in predicate p0) having the same value in a second source operand, e.g., source operand r2, placing the value into shared register ur2 and setting predicate p0 to true for those threads that were just collected.

(50) At this point, a set of threads having uniform source operand values may have been collected (identified by predicate p0), with the common source operand values being stored in shared registers **134**. A uniform instruction uiadd may then be executed by serial execution unit **132** using the uniform source operand values in shared registers ur2 and ur5 and placing the result in shared register ur9. Because the instruction is to be executed by a serial execution unit **132** for at least one thread, no predicate guard is needed. A mov instruction may optionally be performed to move the result to a private register for each thread in the set of threads for which the instruction was executed (e.g., by using predicate p0 as a predicate guard for the instruction).

(51) The set of remaining active threads pending execution may be updated (e.g., to exclude the threads from a next iteration of the loop) using: @p1 mov p1, ~p0

The peeling loop may then proceed to the next iteration with those threads for which the instruction

is still pending execution (e.g., returning to Ltop), ultimately terminating when no threads remaining pending execution (e.g., where branch.u.any evaluates to false or null).

(52) In some embodiments, software compiler logic **150** may further optimize (or attempt to optimize) the peeling loop, e.g., to reduce the number of levels in the peeling loop or flatten the peeling loop entirely. For example, as discussed above, software compiler logic **150** may be able to analyze program code to determine whether an instruction exhibits execution and/or operand uniformity (e.g., by performing iterative flow analysis). Where software compiler logic **150** is able to determine that the instructions exhibit complete operand uniformity, the peeling loop may be eliminated entirely, e.g., generating uniform instructions as discussed herein. Software compiler logic **150**, for example, may be able to determine that the instruction is to be executed by a single thread, based on which software compiler logic **150** may conclude that the instruction exhibits complete operand uniformity (e.g., the single thread is uniform with itself).

(53) In some cases, software compiler logic **150** may be able to determine that an instruction exhibits partial operand uniformity, e.g., where certain operands share the same value across all threads in a group of active threads. In such cases, software compiler logic **150** may copy the operand values to shared registers (as necessary) outside the peeling loop, and partition and unwind the execution of the group of threads based on the remaining source operands. As an illustrative example, with reference to the addition operation previously discussed, software compiler logic **150** may determine that the values of source operand r2 is the same for all threads in an active thread group, in which case software compiler logic **150** may generate the following instructions (with a simplified peeling loop):

(54) TABLE-US-00003 set p1, pTrue // 1 mov ur2, r2 // 2 Ltop: collect p0, ur5, r5, p1 // 3
 uiadd ur9, ur2, ur5 // 4 @p0 mov r9, ur9 // 5 @p1 mov p1, ~p0 // 6 @p1 branch.u.any Ltop // 7

(55) In some cases, software compiler logic **150** may be able to determine that one or more of the source operands is generated by a uniform instruction such that the value need not be copied to a shared register **134** (e.g., because the source operand value is already stored in a shared register **134**). In such cases, software compiler logic **150** may be able to eliminate the mov instruction for the source operand in the above peeling loop (and the collect instruction for the source operand in the peeling loop). Software compiler logic **150**, likewise, may be able to determine that the result of the instruction is only used by other uniform instructions, such that the result need not be copied to a private register **133**. In such cases, the mov instruction may be eliminated from the peeling loop (e.g., the @p0 mov r9, ur9 instruction may be removed from the described peeling loops).

(56) In some cases, software compiler logic **150** may apply branch predication to optimize peeling loops. Software compiler logic **150**, for example, may be able to determine that an instruction is to be executed for multiple branch paths and that the instruction exhibits complete operand uniformity or substantial operand uniformity in each branch (e.g., such that optimization would be beneficial). Rather than generate a collect peeling loop for each branch path to be serially executed, software compiler logic **150** may generate a combined peeling loop in which instructions for each branch path are included within the same peeling loop but are predicated with different conditions (e.g., based on the conditional control construct that resulted in the different branch paths). In some cases, software compiler logic **150** may further optimize the combined peeling loop, e.g., to reduce the number of levels in the peeling loop or flatten the peeling loop entirely, as described herein.

(57) In some embodiments, software compiler logic **150** may operate to promote values from private registers **133** to shared registers **134**, e.g., upon a determination that they are uniform across all threads in a group of active threads, which may facilitate the peeling loop simplifications just described. For example, in some embodiments, analysis of program code by software compiler logic **150** may reveal cohesive regions of the program and/or convergent write operations in the program code, where it may be expected that instruction operands may have uniform values across a group of active threads. Software compiler logic **150** may employ different heuristics to determine when to promote operands within these portions of the program code. For example, in

some embodiments, software compiler logic **150** may consider the number of conversions emitted between private and shared registers, which may provide an indication of whether an instruction can accept a uniform register as an operand or whether a uniform version of an instruction (e.g., a uniform instruction) exists. In some embodiments, software compiler logic **150** may consider register pressure, e.g., on both private and shared register spaces, and may look to balance pressure between the private and shared spaces so as to avoid register spills. In some embodiments, software compiler logic **150** may prioritize reducing pressure in a private register space, as register spills experienced therein may be more expensive. For example, in some embodiments, shared register spills may be placed in private registers, whereas private register spills may go to memory (e.g., local memory **136**).

(58) In some embodiments, software compiler logic **150** may operate to recast instructions to align with operand types permitted by an underlying execution unit (e.g., as supported by parallel execution units **131** or serial execution units **132**). For example, in some embodiments, parallel execution units **131** may support the use of shared register operands for certain operands of an instruction and software compiler logic **150** may operate to reorder the operands of an operation to align any operand type mismatches. By way of example, a parallel execution unit **131** may support an iadd instruction in which shared registers are permitted for the last source operand. Software compiler logic **150** may encounter an addition operation where a shared register is used for the first source operand (but not the second), in which case software compiler logic **150** may recast the operation such that the shared register is provided as the last operand, permitting execution of the iadd instruction on parallel execution unit **131**. Likewise, if a serial execution unit **132** does not support execution of particular operation, software compiler logic **150** may convert any shared register operands (e.g., operands that are to be stored in shared registers **134**) to private register operands (e.g., operands that are to be stored in private registers **133**), such that the operation may be executed using parallel execution unit **131**.

(59) In some embodiments, parallel processor **120** (and its multiprocessors **130**) may not support execution of particular instructions using a parallel execution unit **131**. For instance, because compute logic is replicated for each thread lane in embodiments, parallel execution units **131** may not support instructions that require a significant amount of silicon real estate to implement (e.g., TMA operations). Furthermore, some instructions may be too expensive to execute using a parallel execution unit **132**. For example, certain operations may take a substantial amount of time to execute (e.g., a TMA operation may take several hundreds or thousands of clock cycles to execute) but may only be executed by a few threads at a time, reducing the overall utilization rate of the parallel execution unit **131**. Other operations may not be effectively parallelized and, thus, may not be efficiently executed by a parallel execution unit **131**. Software compiler logic **150** may analyze the program code to determine such instances and may adapt the program code for execution on a serial execution unit **132**.

(60) For example, as discussed above, software compiler logic **150** may generate a peeling loop (e.g., a “collect” peeling loop) that may operate to partition and unwind execution of the group of threads into one or more sets of threads having shared source operands values that are executed once (e.g., for all threads in a particular set of threads) in serial fashion (e.g., for each of the one or more set of threads) using serial execution unit **132**. However, such peeling loops may not be used for certain types of instructions, for example, those that generate results dependent on the number and/or identity of threads executing the instruction (e.g., in-memory reduction operations or counters that count the number of threads executing an instruction). In such cases, software compiler logic **150** may generate a peeling loop that may operate to unwind and execute the instruction one thread at a time on a serial execution unit **132**.

(61) In some embodiments, software compiler logic **150** may make use of a select instruction to select a single thread from a group of active threads for execution. In some embodiments, for example, the select instruction may take the following form: @pg select pd, ps

In some embodiments, the select instruction may operate to choose a single thread at random from a group of active threads pending execution of an instruction, while in other embodiments, a first thread (e.g., having a lowest thread ID) in the group of active threads may be chosen. Source operand *ps* may be a predicate mask identifying a group of active threads from which a selection is made. Destination predicate *pd* may be set to true for the chosen thread, with any other active threads being set to false. In some embodiments, the select instruction may be fully predicated using predicate guard *@pg*.

(62) In some embodiments, for example, software compiler logic **150** may generate a “select” peeling loop, where each iteration of the peeling loop may include a select instruction that selects a single thread from a group of active threads and instructions to copy the source operand values of the selected instruction to shared registers **134** (as necessary). A uniform instruction may then be executed by serial execution unit **132** using the shared source operands, the result of which may, optionally, be copied back to a private register for threads in the set of threads for which the instruction was executed. The source predicate mask identifying active threads may be updated (e.g., to exclude the thread for which the instruction was just executed) and the peeling loop may proceed to the next iteration, ultimately terminating when no threads remain pending execution.

(63) As an illustrative example, software compiler logic **150** may generate the following peeling loop for a reduction operation *red* having two source operands *r0* and *r1*, which may be private registers (e.g., in the register space of each thread), and a destination operand *r2*, which may be a private register (e.g., in the register space of each thread):

(64) TABLE-US-00004 *set p1, pTrue // 1 Ltop: select p0, p1 // 2 @p0 mov ur0, r0 // 3 @p0
mov ur1, r1 // 4 ured ur2, [ur0], ur1 // 5 @p0 mov r2, ur2 // 6 @p1 mov p1, ~p0 // 7 @p1
branch.u.any Ltop // 8*

The predicate *p1* may be used to store the active threads for which the instruction is pending execution. As illustrated herein, predicate *p1* is initialized to true (by *pTrue*), e.g., for all threads in a group of threads (e.g., in a warp), but in other cases predicate *p1* may inherit the predication of the instruction being executed, e.g., where the operation is to be performed for a subset of a group of threads (e.g., a half warp, a quarter warp, etc.).

(65) The loop may begin with a select instruction that may operate to select a thread from the group of threads, setting the predicate *p0* to true for the selected thread (e.g., with all other active threads set to false). The values of source operand *r0* and *r1* for the selected thread may then be copied to respective registers in shared register **134**. A uniform instruction *ured* may then be executed by serial execution unit **132** using the shared source operand values *ur0* and *ur1* in shared registers **134**, placing the result in shared register *ur0*. Because the instruction is to be executed by a serial execution unit **132** for at least one thread, no predicate guard is needed. A *mov* instruction may optionally be performed to move the result back to a private register of the selected thread for which the instruction was executed (e.g., by using predicate *p0* as a predicate guard for the instruction).

(66) The set of remaining active threads pending execution may be updated (e.g., to exclude the thread from a next iteration of the loop) using: *@p1 mov p1, ~p0*

The peeling loop may then proceed to the next iteration with those threads for which the instruction is still pending execution (e.g., returning to *Ltop*), ultimately terminating when no threads remaining pending execution.

(67) In some embodiments, software compiler logic **150** may further optimize the peeling loop, e.g., by utilizing single thread execution analysis to simplify the loop. Software compiler logic **150**, for example, may be able to analyze program code to determine whether an instruction is executed for at least one thread, at most one thread, or exactly one thread. Where software compiler logic **150** is able to determine that an instruction is to be executed by at most one thread, it may not be necessary to perform the select instruction or use a peeling loop. Where software compiler logic **150** is able to determine that an instruction is to be executed by at least one thread execution of the

uniform instruction need not be predicate (e.g., execution of uired need not be predicated). Where both conditions are true, e.g., where software compiler logic **150** is able to determine that exactly one thread executes the instruction, neither the peeling loop nor predication is necessary. For example, with reference to the previous red operation, where software compiler logic **150** is able to determine that the operation is to be performed by exactly one thread, software compiler logic **150** may generate the following instructions:

(68) TABLE-US-00005 mov ur0, r0 mov ur1, r1 uired ur2, [ur0], ur1 mov r2, ur2

(69) Furthermore, as discussed above with respect to the “collect” peeling loop, software compiler logic **150** may be able to determine that one or more of the source operands is generated by a uniform instruction, e.g., the source operand is a shared register **134**, such that the value need not be copied. Software compiler logic **150**, likewise, may be able to determine that the result of the operation is only used by another uniform instruction, such that the result need not be copied back to a private register **133** (e.g., the @p0 mov r2, ur2 instruction may be removed). And in some embodiments, software compiler logic **150** may operate to promote values from private registers **133** to shared registers **134**, e.g., upon a determination that they are uniform across all threads in a group of active threads, which may facilitate these simplifications.

(70) In some embodiments, the single thread execution analysis performed software compiler logic **150** may examine instruction predicates (e.g., instruction predicate guards) to determine whether an instruction is to be executed by at least, at most, or exactly one thread. Software compiler logic **150**, for example, may determine that the select instruction results in selection of exactly one thread, e.g., resulting in a single-thread predicate. As another example, software compiler logic **150** may identify instances where the program code identifies a particular thread within a group of active threads (e.g., a specific lane_id) and uses this identification to define a predicate that is only true for that thread (e.g., comparing the lane_id to an invariant operand mask), e.g., resulting in a single-thread predicate. In some embodiments, software compiler logic **150** may be able to identify single thread blocks (e.g., where a thread block only includes a single block).

(71) In some embodiments, software compiler logic **150** may also track the propagation of single-thread predicates and may be able to draw conclusions about predicates that are derived from a single-thread predicate. Software compiler logic **150**, for example, may determine that instructions predicated by a single-thread predicate, or branch paths executed based on a single-thread predicate are also executed for a single thread. Software compiler logic **150** may also determine that predicates resulting from a logical operation with a single-thread predicate (e.g., AND, OR, SET, NOT) exhibit single thread execution properties. Software compiler logic **150**, for instance, may determine that a predicate resulting from an AND operation with a single-thread predicate indicates that at most one thread is executed or that a predicate resulting from an OR operation with a single-thread predicate indicates that at least one thread is executed.

(72) In some embodiments, computing system **100** may execute binary code **154** generated by software compiler logic **150** using general purpose processor **110** and parallel processors **120**. As discussed above, in compiling binary code **154**, software compiler logic **150** may determine whether to perform particular instructions using parallel execution units **131** or serial execution units **132**. Where a determination is made to perform an instruction on serial execution unit **132**, software compiler logic **150** may have included a peeling loop in binary code **154**. Binary code **154**, for instance, may include one or more collect peeling loops (or optimized versions thereof) that, when executed, may operate to partition a group of threads into sub-groups and perform the instruction once for each sub-group using serial execution unit **132**. Binary code **154** may also include one or more select peeling loops (or optimized versions thereof) that, when executed, may perform the instruction for a group of threads one at a time on serial execution unit **132**.

(73) FIGS. 2 and 3 illustrate example methods in accordance with embodiments of the present disclosure. For the sake of simplicity and clarity, these methods are depicted and described as a series of operations. However, in accordance with the present disclosure, such operations may be

performed in other orders and/or concurrently, and with other operations not presented or described herein. Furthermore, not all illustrated operations may be required in implementing methods in accordance with the present disclosure. Those of skill in the art will also understand and appreciate that the methods could be represented as a series of interrelated states or events via a state diagram. Additionally, it will be appreciated that the disclosed methods are capable of being stored on an article of manufacture. The term “article of manufacture,” as used herein, is intended to encompass a computer-readable device or storage media provided with a computer program and/or executable instructions that, when executed, affect one or more operations.

(74) FIG. 2 illustrates a flow diagram of an example method **200** for executing an instruction on a serial execution unit using a collect peeling loop, according to at least one embodiment. In a SIMT programming and execution model, instructions may be executed by multiple threads concurrently, with each thread reading potentially distinct input data in respective source operands and generating potentially distinct results to respective destination operands. However, it is often the case that multiple threads concurrently execute the same deterministic instruction on uniform sets of input operands to generate the same result. In such cases, the redundant execution of the instruction by each thread using a parallel execution unit (e.g., parallel execution units **131**) unnecessarily increases power consumption and resource utilization. A peeling loop may be used to partition and unwind execution of the group of threads into one or more sets of threads having shared source operand values that are executed once (e.g., for all threads in a particular set of threads) in serial fashion (e.g., for each of the one or more set of threads).

(75) At operation **210**, processing logic may define a first predicate mask that may identify a group of active threads for which an instruction is pending execution. In some cases, the predicate mask may be initialized to true for all threads in a group of threads. In others, the predicate mask may be initialized with an inherited predicate mask of the instruction being executed.

(76) At operation **220**, processing logic may begin the peeling loop by collecting a set of threads having uniform source operand values (e.g., where each source operand of the instruction has the same value across all threads) and (as necessary) place the source operand values in shared registers (e.g., shared registers **134**) accessible by a serial execution unit (e.g., serial execution unit **132**).

(77) At block **221**, processing logic may collect a set of threads having the same value in an initial source operand, placing the operand value into a shared register (e.g., into shared registers **134**) and setting a loop predicate mask to true for those threads that were collected (e.g., with all other active threads set to false). In some embodiments, the selection of an anchor source operand value (e.g., based on which the collection is performed) may be identified at random from amongst the threads being collected, while in other embodiments, the source operand value of a first thread (e.g., having a lowest thread ID) may be used as the anchor value for the collect operation.

(78) At block **222**, processing logic may repeat the collect operation recursively for the remaining source operands, e.g., by collecting a subset of threads from those previously collected (e.g., from the threads in the loop predicate mask) having the same value in another source operand, placing the operand value into a shared register (e.g., into shared registers **134**) and updating the loop predicate mask to be true for the subset of threads just collected. As before, in some embodiments, the selection of an anchor source operand value may be identified at random, while in others, the source operand value of a first thread may be used as the anchor value for the collect operation.

(79) At operation **230**, once the collect operation has been performed for (e.g., all) source operands and a set of threads having uniform source operand values has been collected (e.g., identified by the loop predicate mask), the processing logic may execute a uniform instruction using the uniform source operand values on a serial execution unit (e.g., serial execution unit **132**), with a result being placed into a shared register (e.g., into shared registers **134**). Because the instruction is to be executed by a serial execution unit **132** for at least one thread, no predicate guard is needed.

(80) At operation **240**, the processing logic may optionally copy the result back to a private register

(e.g., private registers **133**) for each collected thread in the set of threads for which the instruction was executed, for example, by using the loop predicate mask as a predicate guard for the copy operation.

(81) At operation **250**, processing logic may update the first predicate mask to identify those threads for which execution of the instruction remains pending (e.g., to exclude the threads collected in the current iteration from being considered in a next iteration of the loop).

(82) At operation **260**, processing logic may determine whether any threads remain pending execution. If so, processing logic may return to operation **220**, otherwise the loop may terminate and processing logic may proceed with a next instruction.

(83) FIG. **3** illustrates a flow diagram of an example method **300** for executing an instruction on a serial execution unit using a select peeling loop, according to at least one embodiment. In some cases, a parallel processor (e.g., parallel processor **120**) may not support execution of an instruction using a parallel execution unit (e.g., parallel execution units **131** of multiprocessors **130**) or an instruction may be too expensive to execute using a parallel execution unit. Some instructions may also generate results that depend on the number and/or identity of threads executing the instruction for which a collect peeling loop may not be employed. In such cases, a select peeling loop may be used to unwind and execute the instruction one thread at a time on a serial execution unit (e.g., serial execution unit **132**).

(84) At operation **310**, processing logic may define a first predicate mask that may identify a group of active threads for which an instruction is pending execution. In some cases, the predicate mask may be initialized to true for all threads in the group of threads. In others, the predicate mask may be initialized with an inherited predicate mask of the instruction being executed.

(85) At operation **320**, processing logic may begin the peeling loop by selecting a thread from the group of active threads and setting a loop predicate mask to true for that thread (e.g., with all other active threads set to false). In some embodiments, the selection of the thread may be performed at random, while in others, the first thread in the group of threads (e.g., having a lowest thread ID) may be selected.

(86) At operation **330**, processing logic may copy (as necessary) the values of private source operands (e.g., in private registers **133**) to shared registers (e.g., to shared registers **134**) to obtain shared source operands.

(87) At operation **340**, the processing logic may execute a uniform instruction using the shared source operands on a serial execution unit (e.g., serial execution unit **132**), with a result being placed into a shared register (e.g., into shared registers **134**). Because the instruction is to be executed by a serial execution unit **132** for at least one thread, no predicate guard is needed.

(88) At operation **350**, the processing logic may optionally copy the result back to a private register (e.g., private registers **133**) for the selected thread for which the instruction was executed (e.g., using the loop predicate mask as a predicate guard for the copy operation).

(89) At operation **360**, processing logic may update the first predicate mask to identify those threads for which execution of the instruction remains pending (e.g., to exclude the thread selected in the current iteration from being considered in a next iteration of the loop).

(90) At operation **370**, processing logic may determine whether any threads remain pending execution. If so, processing logic may return to operation **320**, otherwise the loop may terminate and processing logic may proceed with a next instruction.

(91) FIG. **4** illustrates a computer system **400**, according to at least one embodiment. In at least one embodiment, computer system **400** is configured to implement various processes and methods described throughout this disclosure.

(92) In at least one embodiment, computer system **400** comprises at least one central processing unit ("CPU") **402** that is connected to a communication bus **410** implemented using any suitable protocol, such as PCI ("Peripheral Component Interconnect"), peripheral component interconnect express ("PCI-Express"), AGP ("Accelerated Graphics Port"), HyperTransport, or any other bus or

point-to-point communication protocol(s). In at least one embodiment, computer system **400** includes a main memory **404**, which may take form of random-access memory (“RAM”). Control logic (e.g., implemented as hardware, software, or a combination thereof) and data are stored in main memory **404**. In at least one embodiment, a network interface subsystem (“network interface”) **422** provides an interface to other computing devices and networks for receiving data from and transmitting data to other systems with computer system **400**.

(93) In at least one embodiment, computer system **400** includes one or more input devices **408**, a parallel processing system **412**, and one or more display devices **406** that can be implemented using a conventional cathode ray tube (“CRT”), a liquid crystal display (“LCD”), a light emitting diode (“LED”) display, a plasma display, or other suitable display technologies. In at least one embodiment, user input is received from input devices **408** such as keyboard, mouse, touchpad, microphone, etc. In at least one embodiment, each module described herein can be situated on a single semiconductor platform to form a processing system.

(94) Software compiler logic **150** may be used to perform software compilation operations associated with one or more embodiments, including generating and optimizing peeling loops to execute instructions on serial execution units. Details regarding software compiler logic **150** are provided herein in conjunction with FIG. **1**. In at least one embodiment, software compiler logic **150** and the program code it may generate may be used in the system of FIG. **4** for performing software compilation operations, including generating and optimizing peeling loops to execute instructions on serial execution units, and executing instructions on serial execution units.

(95) In at least one embodiment, computer programs in form of machine-readable executable code or computer control logic algorithms are stored in main memory **404** and/or secondary storage. Computer programs, if executed by one or more processors, enable system **400** to perform various functions in accordance with at least one embodiment. In at least one embodiment, memory **404**, storage, and/or any other storage are possible examples of computer-readable media. In at least one embodiment, secondary storage may refer to any suitable storage device or system such as a hard disk drive and/or a removable storage drive, representing a floppy disk drive, a magnetic tape drive, a compact disk drive, digital versatile disk (“DVD”) drive, recording device, universal serial bus (“USB”) flash memory, etc.

(96) In at least one embodiment, architecture and/or functionality of various previous figures are implemented in the context of CPU **402**, parallel processing system **412**, an integrated circuit capable of at least a portion of capabilities of both CPU **402** and parallel processing system **412**, a chipset (e.g., a group of integrated circuits designed to work and sold as a unit for performing related functions, etc.), and/or any suitable combination of integrated circuit(s). In at least one embodiment, architecture and/or functionality of various previous figures are implemented in the context of a general computer system, a circuit board system, a game console system dedicated for entertainment purposes, an application-specific system, and more. In at least one embodiment, computer system **400** may take the form of a desktop computer, a laptop computer, a tablet computer, a server, a supercomputer, a smart-phone (e.g., a wireless, hand-held device), a personal digital assistant (“PDA”), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, a mobile phone device, a television, a workstation, a game console, an embedded system, and/or any other type of logic device.

(97) In at least one embodiment, parallel processing system **412** includes a plurality of parallel processing units (“PPUs”) **414** and associated memories **416**. In at least one embodiment, PPUs **414** are connected to a host processor or other peripheral devices via an interconnect **418** and a switch **420** or multiplexer. In at least one embodiment, parallel processing system **412** distributes computational tasks across PPUs **414** which can be parallelizable for example, as part of distribution of computational tasks across multiple graphics processing unit (“GPU”) thread blocks. In at least one embodiment, memory is shared and accessible (e.g., for read and/or write access) across some or all of PPUs **414**, although such shared memory may incur performance penalties

relative to use of local memory and registers resident to a PPU **414**. In at least one embodiment, operation of PPUs **414** is synchronized through use of a command such as `_syncthreads()`, wherein all threads in a block (e.g., executed across multiple PPUs **414**) to reach a certain point of execution of code before proceeding.

(98) FIG. 5 illustrates a parallel processing unit (“PPU”) **500**, according to at least one embodiment. In at least one embodiment, PPU **500** is configured with machine-readable code that, if executed by PPU **500**, causes PPU **500** to perform some or all of processes and techniques described throughout this disclosure. In at least one embodiment, PPU **500** is a multi-threaded processor that is implemented on one or more integrated circuit devices and that utilizes multithreading as a latency-hiding technique designed to process computer-readable instructions (also referred to as machine-readable instructions or simply instructions) on multiple threads in parallel. In at least one embodiment, a thread refers to a thread of execution and is an instantiation of a set of instructions configured to be executed by PPU **500**. In at least one embodiment, PPU **500** is a graphics processing unit (“GPU”) configured to implement a graphics rendering pipeline for processing three-dimensional (“3D”) graphics data in order to generate two-dimensional (“2D”) image data for display on a display device such as a liquid crystal display (“LCD”) device. In at least one embodiment, PPU **500** is utilized to perform computations such as linear algebra operations and machine-learning operations. FIG. 5 illustrates an example parallel processor for illustrative purposes only and should be construed as a non-limiting example of processor architectures contemplated within scope of this disclosure and that any suitable processor may be employed to supplement and/or substitute for same.

(99) In at least one embodiment, one or more PPUs **500** are configured to accelerate High Performance Computing (“HPC”), data center, and machine learning applications. In at least one embodiment, PPU **500** is configured to accelerate deep learning systems and applications including following non-limiting examples: autonomous vehicle platforms, deep learning, high-accuracy speech, image, text recognition systems, intelligent video analytics, molecular simulations, drug discovery, disease diagnosis, weather forecasting, big data analytics, astronomy, molecular dynamics simulation, financial modeling, robotics, factory automation, real-time language translation, online search optimizations, and personalized user recommendations, and more.

(100) In at least one embodiment, PPU **500** includes an Input/Output (“I/O”) unit **506**, a front-end unit **510**, a scheduler unit **512**, a work distribution unit **514**, a hub **516**, a crossbar (“XBar”) **520**, one or more general processing clusters (“GPCs”) **518**, and one or more partition units (“memory partition units”) **522**. In at least one embodiment, PPU **500** is connected to a host processor or other PPUs **500** via one or more high-speed GPU interconnects (“GPU interconnects”) **508**. In at least one embodiment, PPU **500** is connected to a host processor or other peripheral devices via a system bus **502**. In at least one embodiment, PPU **500** is connected to a local memory comprising one or more memory devices (“memory”) **504**. In at least one embodiment, memory devices **504** include one or more dynamic random access memory (“DRAM”) devices. In at least one embodiment, one or more DRAM devices are configured and/or configurable as high-bandwidth memory (“HBM”) subsystems, with multiple DRAM dies stacked within each device.

(101) In at least one embodiment, high-speed GPU interconnect **508** may refer to a wire-based multi-lane communications link that is used by systems to scale and include one or more PPUs **500** combined with one or more central processing units (“CPUs”), supports cache coherence between PPUs **500** and CPUs, and CPU mastering. In at least one embodiment, data and/or commands are transmitted by high-speed GPU interconnect **508** through hub **516** to/from other units of PPU **500** such as one or more copy engines, video encoders, video decoders, power management units, and other components which may not be explicitly illustrated in FIG. 5.

(102) In at least one embodiment, I/O unit **506** is configured to transmit and receive communications (e.g., commands, data) from a host processor (not illustrated in FIG. 5) over system bus **502**. In at least one embodiment, I/O unit **506** communicates with host processor

directly via system bus **502** or through one or more intermediate devices such as a memory bridge. In at least one embodiment, I/O unit **506** may communicate with one or more other processors, such as one or more of PPUs **500** via system bus **502**. In at least one embodiment, I/O unit **506** implements a Peripheral Component Interconnect Express (“PCIe”) interface for communications over a PCIe bus. In at least one embodiment, I/O unit **506** implements interfaces for communicating with external devices.

(103) In at least one embodiment, I/O unit **506** decodes packets received via system bus **502**. In at least one embodiment, at least some packets represent commands configured to cause PPU **500** to perform various operations. In at least one embodiment, I/O unit **506** transmits decoded commands to various other units of PPU **500** as specified by commands. In at least one embodiment, commands are transmitted to front-end unit **510** and/or transmitted to hub **516** or other units of PPU **500** such as one or more copy engines, a video encoder, a video decoder, a power management unit, etc. (not explicitly illustrated in FIG. 5). In at least one embodiment, I/O unit **506** is configured to route communications between and among various logical units of PPU **500**.

(104) In at least one embodiment, a program executed by host processor encodes a command stream in a buffer that provides workloads to PPU **500** for processing. In at least one embodiment, a workload comprises instructions and data to be processed by those instructions. In at least one embodiment, a buffer is a region in a memory that is accessible (e.g., read/write) by both a host processor and PPU **500**—a host interface unit may be configured to access that buffer in a system memory connected to system bus **502** via memory requests transmitted over system bus **502** by I/O unit **506**. In at least one embodiment, a host processor writes a command stream to a buffer and then transmits a pointer to a start of a command stream to PPU **500** such that front-end unit **510** receives pointers to one or more command streams and manages one or more command streams, reading commands from command streams and forwarding commands to various units of PPU **500**.

(105) In at least one embodiment, front-end unit **510** is coupled to scheduler unit **512** that configures various GPCs **518** to process tasks defined by one or more command streams. In at least one embodiment, scheduler unit **512** is configured to track state information related to various tasks managed by scheduler unit **512** where state information may indicate which of GPCs **518** a task is assigned to, whether task is active or inactive, a priority level associated with task, and so forth. In at least one embodiment, scheduler unit **512** manages execution of a plurality of tasks on one or more of GPCs **518**.

(106) In at least one embodiment, scheduler unit **512** is coupled to work distribution unit **514** that is configured to dispatch tasks for execution on GPCs **518**. In at least one embodiment, work distribution unit **514** tracks a number of scheduled tasks received from scheduler unit **512** and work distribution unit **514** manages a pending task pool and an active task pool for each of GPCs **518**. In at least one embodiment, pending task pool comprises a number of slots (e.g., 32 slots) that contain tasks assigned to be processed by a particular GPC **518**; an active task pool may comprise a number of slots (e.g., 4 slots) for tasks that are actively being processed by GPCs **518** such that as one of GPCs **518** completes execution of a task, that task is evicted from that active task pool for GPC **518** and another task from a pending task pool is selected and scheduled for execution on GPC **518**. In at least one embodiment, if an active task is idle on GPC **518**, such as while waiting for a data dependency to be resolved, then that active task is evicted from GPC **518** and returned to that pending task pool while another task in that pending task pool is selected and scheduled for execution on GPC **518**.

(107) In at least one embodiment, work distribution unit **514** communicates with one or more GPCs **518** via XBar **520**. In at least one embodiment, XBar **520** is an interconnect network that couples many of units of PPU **500** to other units of PPU **500** and can be configured to couple work distribution unit **514** to a particular GPC **518**. In at least one embodiment, one or more other units of PPU **500** may also be connected to XBar **520** via hub **516**.

(108) In at least one embodiment, tasks are managed by scheduler unit **512** and dispatched to one

of GPCs **518** by work distribution unit **514**. In at least one embodiment, GPC **518** is configured to process task and generate results. In at least one embodiment, results may be consumed by other tasks within GPC **518**, routed to a different GPC **518** via XBar **520**, or stored in memory **504**. In at least one embodiment, results can be written to memory **504** via partition units **522**, which implement a memory interface for reading and writing data to/from memory **504**. In at least one embodiment, results can be transmitted to another PPU **500** or CPU via high-speed GPU interconnect **508**. In at least one embodiment, PPU **500** includes a number U of partition units **522** that is equal to a number of separate and distinct memory devices **504** coupled to PPU **500**, as described in more detail herein in conjunction with FIG. **28**.

(109) In at least one embodiment, a host processor executes a driver kernel that implements an application programming interface (“API”) that enables one or more applications executing on a host processor to schedule operations for execution on PPU **500**. In at least one embodiment, multiple compute applications are simultaneously executed by PPU **500** and PPU **500** provides isolation, quality of service (“QoS”), and independent address spaces for multiple compute applications. In at least one embodiment, an application generates instructions (e.g., in form of API calls) that cause a driver kernel to generate one or more tasks for execution by PPU **500** and that driver kernel outputs tasks to one or more streams being processed by PPU **500**. In at least one embodiment, each task comprises one or more groups of related threads, which may be referred to as a warp. In at least one embodiment, a warp comprises a plurality of related threads (e.g., 32 threads) that can be executed in parallel. In at least one embodiment, cooperating threads can refer to a plurality of threads including instructions to perform task and that exchange data through shared memory. In at least one embodiment, threads and cooperating threads are described in more detail in conjunction with FIG. **28**.

(110) Software compiler logic **150** may be used to perform software compilation operations associated with one or more embodiments, including generating and optimizing peeling loops to execute instructions on serial execution units. Details regarding software compiler logic **150** are provided herein in conjunction with FIG. **1**. In at least one embodiment, software compiler logic **150** and the program code it may generate may be used in the system of FIG. **5** for performing software compilation operations, including generating and optimizing peeling loops to execute instructions on serial execution units, and executing instructions on serial execution units.

(111) FIG. **6** illustrates a general processing cluster (“GPC”) **600**, according to at least one embodiment. In at least one embodiment, GPC **600** is GPC **518** of FIG. **5**. In at least one embodiment, each GPC **600** includes a number of hardware units for processing tasks and each GPC **600** includes a pipeline manager **602**, a pre-raster operations unit (“preROP”) **604**, a raster engine **608**, a work distribution crossbar (“WDX”) **616**, a memory management unit (“MMU”) **618**, one or more Data Processing Clusters (“DPCs”) **606**, and any suitable combination of parts.

(112) In at least one embodiment, operation of GPC **600** is controlled by pipeline manager **602**. In at least one embodiment, pipeline manager **602** manages configuration of one or more DPCs **606** for processing tasks allocated to GPC **600**. In at least one embodiment, pipeline manager **602** configures at least one of one or more DPCs **606** to implement at least a portion of a graphics rendering pipeline. In at least one embodiment, DPC **606** is configured to execute a vertex shader program on a programmable streaming multi-processor (“SM”) **614**. In at least one embodiment, pipeline manager **602** is configured to route packets received from a work distribution unit to appropriate logical units within GPC **600**, in at least one embodiment, and some packets may be routed to fixed function hardware units in preROP **604** and/or raster engine **608** while other packets may be routed to DPCs **606** for processing by a primitive engine **612** or SM **614**. In at least one embodiment, pipeline manager **602** configures at least one of DPCs **606** to implement a neural network model and/or a computing pipeline.

(113) In at least one embodiment, preROP unit **604** is configured, in at least one embodiment, to route data generated by raster engine **608** and DPCs **606** to a Raster Operations (“ROP”) unit in

partition unit **522**, described in more detail above in conjunction with FIG. 5. In at least one embodiment, preROP unit **604** is configured to perform optimizations for color blending, organize pixel data, perform address translations, and more. In at least one embodiment, raster engine **608** includes a number of fixed function hardware units configured to perform various raster operations, in at least one embodiment, and raster engine **608** includes a setup engine, a coarse raster engine, a culling engine, a clipping engine, a fine raster engine, a tile coalescing engine, and any suitable combination thereof. In at least one embodiment, setup engine receives transformed vertices and generates plane equations associated with geometric primitive defined by vertices; plane equations are transmitted to a coarse raster engine to generate coverage information (e.g., an x, y coverage mask for a tile) for primitive; output of a coarse raster engine is transmitted to a culling engine where fragments associated with a primitive that fail a z-test are culled, and transmitted to a clipping engine where fragments lying outside a viewing frustum are clipped. In at least one embodiment, fragments that survive clipping and culling are passed to a fine raster engine to generate attributes for pixel fragments based on plane equations generated by a setup engine. In at least one embodiment, an output of raster engine **608** comprises fragments to be processed by any suitable entity, such as by a fragment shader implemented within DPC **606**.

(114) In at least one embodiment, each DPC **606** included in GPC **600** comprises an M-Pipe Controller (“MPC”) **610**; primitive engine **612**; one or more SMs **614**; and any suitable combination thereof. In at least one embodiment, MPC **610** controls operation of DPC **606**, routing packets received from pipeline manager **602** to appropriate units in DPC **606**. In at least one embodiment, packets associated with a vertex are routed to primitive engine **612**, which is configured to fetch vertex attributes associated with a vertex from memory; in contrast, packets associated with a shader program may be transmitted to SM **614**.

(115) In at least one embodiment, SM **614** comprises a programmable streaming processor that is configured to process tasks represented by a number of threads. In at least one embodiment, SM **614** is multi-threaded and configured to execute a plurality of threads (e.g., 32 threads) from a particular group of threads concurrently and implements a Single-Instruction, Multiple-Data (“SIMD”) architecture where each thread in a group of threads (e.g., a warp) is configured to process a different set of data based on same set of instructions. In at least one embodiment, all threads in group of threads execute a common set of instructions. In at least one embodiment, SM **614** implements a Single-Instruction, Multiple Thread (“SIMT”) architecture wherein each thread in a group of threads is configured to process a different set of data based on that common set of instructions, but where individual threads in a group of threads are allowed to diverge during execution. In at least one embodiment, a program counter, call stack, and execution state is maintained for each warp, enabling concurrency between warps and serial execution within warps when threads within a warp diverge. In another embodiment, a program counter, call stack, and execution state is maintained for each individual thread, enabling equal concurrency between all threads, within and between warps. In at least one embodiment, execution state is maintained for each individual thread and threads executing common instructions may be converged and executed in parallel for better efficiency. At least one embodiment of SM **614** is described in more detail herein.

(116) In at least one embodiment, MMU **618** provides an interface between GPC **600** and a memory partition unit (e.g., partition unit **522** of FIG. 5) and MMU **618** provides translation of virtual addresses into physical addresses, memory protection, and arbitration of memory requests. In at least one embodiment, MMU **618** provides one or more translation lookaside buffers (“TLBs”) for performing translation of virtual addresses into physical addresses in memory.

(117) Software compiler logic **150** may be used to perform software compilation operations associated with one or more embodiments, including generating and optimizing peeling loops to execute instructions on serial execution units. Details regarding software compiler logic **150** are provided herein in conjunction with FIG. 1. In at least one embodiment, software compiler logic

150 and the program code it may generate may be used in the system of FIG. 6 for performing software compilation operations, including generating and optimizing peeling loops to execute instructions on serial execution units, and executing instructions on serial execution units.

(118) FIG. 7 illustrates a memory partition unit **700** of a parallel processing unit (“PPU”), in accordance with at least one embodiment. In at least one embodiment, memory partition unit **700** includes a Raster Operations (“ROP”) unit **702**, a level two (“L2”) cache **704**, a memory interface **706**, and any suitable combination thereof. In at least one embodiment, memory interface **706** is coupled to memory. In at least one embodiment, memory interface **706** may implement 32, 64, 17, 1024-bit data buses, or like, for high-speed data transfer. In at least one embodiment, PPU incorporates U memory interfaces **706** where U is a positive integer, with one memory interface **706** per pair of partition units **700**, where each pair of partition units **700** is connected to a corresponding memory device. For example, in at least one embodiment, PPU may be connected to up to Y memory devices, such as high bandwidth memory stacks or graphics double-data-rate, version 5, synchronous dynamic random access memory (“GDDR5 SDRAM”).

(119) In at least one embodiment, memory interface **706** implements a high bandwidth memory second generation (“HBM2”) memory interface and Y equals half of U. In at least one embodiment, HBM2 memory stacks are located on a physical package with a PPU, providing substantial power and area savings compared with conventional GDDR5 SDRAM systems. In at least one embodiment, each HBM2 stack includes four memory dies with Y=4, with each HBM2 stack including two 128-bit channels per die for a total of 8 channels and a data bus width of 1024 bits. In at least one embodiment, that memory supports Single-Error Correcting Double-Error Detecting (“SECCDED”) Error Correction Code (“ECC”) to protect data. In at least one embodiment, ECC can provide higher reliability for compute applications that are sensitive to data corruption.

(120) In at least one embodiment, PPU implements a multi-level memory hierarchy. In at least one embodiment, memory partition unit **700** supports a unified memory to provide a single unified virtual address space for central processing unit (“CPU”) and PPU memory, enabling data sharing between virtual memory systems. In at least one embodiment frequency of accesses by a PPU to a memory located on other processors is traced to ensure that memory pages are moved to physical memory of PPU that is accessing pages more frequently. In at least one embodiment, high-speed GPU interconnect **508** supports address translation services allowing PPU to directly access a CPU's page tables and providing full access to CPU memory by a PPU.

(121) In at least one embodiment, copy engines transfer data between multiple PPUs or between PPUs and CPUs. In at least one embodiment, copy engines can generate page faults for addresses that are not mapped into page tables and memory partition unit **700** then services page faults, mapping addresses into page table, after which copy engine performs a transfer. In at least one embodiment, memory is pinned (e.g., non-pageable) for multiple copy engine operations between multiple processors, substantially reducing available memory. In at least one embodiment, with hardware page faulting, addresses can be passed to copy engines without regard as to whether memory pages are resident, and a copy process is transparent.

(122) Data from memory **504** of FIG. 5 or other system memory is fetched by memory partition unit **700** and stored in L2 cache **704**, which is located on-chip and is shared between various GPCs, in accordance with at least one embodiment. Each memory partition unit **700**, in at least one embodiment, includes at least a portion of L2 cache associated with a corresponding memory device. In at least one embodiment, lower level caches are implemented in various units within GPCs. In at least one embodiment, each of SMs **614** in FIG. 6 may implement a Level 1 (“L1”) cache wherein that L1 cache is private memory that is dedicated to a particular SM **614** and data from L2 cache **704** is fetched and stored in each L1 cache for processing in functional units of SMs **614**. In at least one embodiment, L2 cache **704** is coupled to memory interface **706** and XBar **520** shown in FIG. 5.

(123) ROP unit **702** performs graphics raster operations related to pixel color, such as color compression, pixel blending, and more, in at least one embodiment. ROP unit **702**, in at least one embodiment, implements depth testing in conjunction with raster engine **608**, receiving a depth for a sample location associated with a pixel fragment from a culling engine of raster engine **608**. In at least one embodiment, depth is tested against a corresponding depth in a depth buffer for a sample location associated with a fragment. In at least one embodiment, if that fragment passes that depth test for that sample location, then ROP unit **702** updates depth buffer and transmits a result of that depth test to raster engine **608**. It will be appreciated that a number of partition units **700** may be different than a number of GPCs and, therefore, each ROP unit **702** can, in at least one embodiment, be coupled to each GPC. In at least one embodiment, ROP unit **702** tracks packets received from different GPCs and determines whether a result generated by ROP unit **702** is to be routed to through XBar **520**.

(124) FIG. **8** illustrates a streaming multi-processor (“SM”) **800**, according to at least one embodiment. In at least one embodiment, SM **800** is SM of FIG. **6**. In at least one embodiment, SM **800** includes an instruction cache **802**, one or more scheduler units **804**, a register **808**, one or more processing cores (“cores”) **810**, one or more special function units (“SFUs”) **812**, one or more load/store units (“LSUs”) **814**, an interconnect network **816**, a shared memory/level one (“L1”) cache **818**, and/or any suitable combination thereof.

(125) In at least one embodiment, a work distribution unit dispatches tasks for execution on general processing clusters (“GPCs”) of parallel processing units (“PPUs”) and each task is allocated to a particular Data Processing Cluster (“DPC”) within a GPC and, if a task is associated with a shader program, that task is allocated to one of SMs **800**. In at least one embodiment, scheduler unit **804** receives tasks from a work distribution unit and manages instruction scheduling for one or more thread blocks assigned to SM **800**. In at least one embodiment, scheduler unit **804** schedules thread blocks for execution as warps of parallel threads, wherein each thread block is allocated at least one warp. In at least one embodiment, each warp executes threads. In at least one embodiment, scheduler unit **804** manages a plurality of different thread blocks, allocating warps to different thread blocks and then dispatching instructions from plurality of different cooperative groups to various functional units (e.g., processing cores **810**, SFUs **812**, and LSUs **814**) during each clock cycle.

(126) In at least one embodiment, Cooperative Groups may refer to a programming model for organizing groups of communicating threads that allows developers to express granularity at which threads are communicating, enabling expression of richer, more efficient parallel decompositions. In at least one embodiment, cooperative launch APIs support synchronization amongst thread blocks for execution of parallel algorithms. In at least one embodiment, applications of conventional programming models provide a single, simple construct for synchronizing cooperating threads: a barrier across all threads of a thread block (e.g., `syncthreads( )` function). However, in at least one embodiment, programmers may define groups of threads at smaller than thread block granularities and synchronize within defined groups to enable greater performance, design flexibility, and software reuse in form of collective group-wide function interfaces. In at least one embodiment, Cooperative Groups enables programmers to define groups of threads explicitly at sub-block (e.g., as small as a single thread) and multi-block granularities, and to perform collective operations such as synchronization on threads in a cooperative group. In at least one embodiment, that programming model supports clean composition across software boundaries, so that libraries and utility functions can synchronize safely within their local context without having to make assumptions about convergence. In at least one embodiment, Cooperative Groups primitives enable new patterns of cooperative parallelism, including producer-consumer parallelism, opportunistic parallelism, and global synchronization across an entire grid of thread blocks.

(127) In at least one embodiment, a dispatch unit **806** is configured to transmit instructions to one

or more functional units and scheduler unit **804** and includes two dispatch units **806** that enable two different instructions from a common warp to be dispatched during each clock cycle. In at least one embodiment, each scheduler unit **804** includes a single dispatch unit **806** or additional dispatch units **806**.

(128) In at least one embodiment, each SM **800**, in at least one embodiment, includes register **808** that provides a set of registers for functional units of SM **800**. In at least one embodiment, register **808** is divided between each functional unit such that each functional unit is allocated a dedicated portion of register **808**. In at least one embodiment, register **808** is divided between different warps being executed by SM **800** and register **808** provides temporary storage for operands connected to data paths of functional units. In at least one embodiment, each SM **800** comprises a plurality of L processing cores **810**, where L is a positive integer. In at least one embodiment, SM **800** includes a large number (e.g., 128 or more) of distinct processing cores **810**. In at least one embodiment, each processing core **810** includes a fully-pipelined, single-precision, double-precision, and/or mixed precision processing unit that includes a floating point arithmetic logic unit and an integer arithmetic logic unit. In at least one embodiment, floating point arithmetic logic units implement IEEE 754-2008 standard for floating point arithmetic. In at least one embodiment, processing cores **810** include 64 single-precision (32-bit) floating point cores, 64 integer cores, 32 double-precision (64-bit) floating point cores, and 8 tensor cores.

(129) Tensor cores are configured to perform matrix operations in accordance with at least one embodiment. In at least one embodiment, one or more tensor cores are included in processing cores **810**. In at least one embodiment, tensor cores are configured to perform deep learning matrix arithmetic, such as convolution operations for neural network training and inferencing. In at least one embodiment, each tensor core operates on a 4×4 matrix and performs a matrix multiply and accumulate operation, $D=A \times B+C$, where A, B, C, and D are 4×4 matrices.

(130) In at least one embodiment, matrix multiply inputs A and B are 16-bit floating point matrices and accumulation matrices C and D are 16-bit floating point or 32-bit floating point matrices. In at least one embodiment, tensor cores operate on 16-bit floating point input data with 32-bit floating point accumulation. In at least one embodiment, 16-bit floating point multiply uses 64 operations and results in a full precision product that is then accumulated using 32-bit floating point addition with other intermediate products for a $4 \times 4 \times 4$ matrix multiply. Tensor cores are used to perform much larger two-dimensional or higher dimensional matrix operations, built up from these smaller elements, in at least one embodiment. In at least one embodiment, an API, such as a CUDA 9 C++ API, exposes specialized matrix load, matrix multiply and accumulate, and matrix store operations to efficiently use tensor cores from a CUDA-C++ program. In at least one embodiment, at a CUDA level, a warp-level interface assumes 16×16 size matrices spanning all 32 threads of warp.

(131) In at least one embodiment, each SM **800** comprises M SFUs **812** that perform special functions (e.g., attribute evaluation, reciprocal square root, and like). In at least one embodiment, SFUs **812** include a tree traversal unit configured to traverse a hierarchical tree data structure. In at least one embodiment, SFUs **812** include a texture unit configured to perform texture map filtering operations. In at least one embodiment, texture units are configured to load texture maps (e.g., a 2D array of texels) from memory and sample texture maps to produce sampled texture values for use in shader programs executed by SM **800**. In at least one embodiment, texture maps are stored in shared memory/L1 cache **818**. In at least one embodiment, texture units implement texture operations such as filtering operations using mip-maps (e.g., texture maps of varying levels of detail), in accordance with at least one embodiment. In at least one embodiment, each SM **800** includes two texture units.

(132) Each SM **800** comprises N LSUs **814** that implement load and store operations between shared memory/L1 cache **818** and register **808**, in at least one embodiment. Interconnect network **816** connects each functional unit to register **808** and LSU **814** to register **808** and shared memory/L1 cache **818** in at least one embodiment. In at least one embodiment, interconnect

network **816** is a crossbar that can be configured to connect any functional units to any registers in register **808** and connect LSUs **814** to register **808** and memory locations in shared memory/L1 cache **818**.

(133) In at least one embodiment, shared memory/L1 cache **818** is an array of on-chip memory that allows for data storage and communication between SM **800** and primitive engine and between threads in SM **800**, in at least one embodiment. In at least one embodiment, shared memory/L1 cache **818** comprises 128 KB of storage capacity and is in a path from SM **800** to a partition unit. In at least one embodiment, shared memory/L1 cache **818**, in at least one embodiment, is used to cache reads and writes. In at least one embodiment, one or more of shared memory/L1 cache **818**, L2 cache, and memory are backing stores.

(134) Combining data cache and shared memory functionality into a single memory block provides improved performance for both types of memory accesses, in at least one embodiment. In at least one embodiment, capacity is used or is usable as a cache by programs that do not use shared memory, such as if shared memory is configured to use half of a capacity, and texture and load/store operations can use remaining capacity. Integration within shared memory/L1 cache **818** enables shared memory/L1 cache **818** to function as a high-throughput conduit for streaming data while simultaneously providing high-bandwidth and low-latency access to frequently reused data, in accordance with at least one embodiment. In at least one embodiment, when configured for general purpose parallel computation, a simpler configuration can be used compared with graphics processing. In at least one embodiment, fixed function graphics processing units are bypassed, creating a much simpler programming model. In a general purpose parallel computation configuration, a work distribution unit assigns and distributes blocks of threads directly to DPCs, in at least one embodiment. In at least one embodiment, threads in a block execute a common program, using a unique thread ID in calculation to ensure each thread generates unique results, using SM **800** to execute program and perform calculations, shared memory/L1 cache **818** to communicate between threads, and LSU **814** to read and write global memory through shared memory/L1 cache **818** and memory partition unit. In at least one embodiment, when configured for general purpose parallel computation, SM **800** writes commands that scheduler unit **804** can use to launch new work on DPCs.

(135) In at least one embodiment, a PPU is included in or coupled to a desktop computer, a laptop computer, a tablet computer, servers, supercomputers, a smart-phone (e.g., a wireless, hand-held device), personal digital assistant (“PDA”), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, and more. In at least one embodiment, a PPU is embodied on a single semiconductor substrate. In at least one embodiment, a PPU is included in a system-on-a-chip (“SoC”) along with one or more other devices such as additional PPUs, memory, a reduced instruction set computer (“RISC”) CPU, a memory management unit (“MMU”), a digital-to-analog converter (“DAC”), and like.

(136) In at least one embodiment, a PPU may be included on a graphics card that includes one or more memory devices. In at least one embodiment, that graphics card may be configured to interface with a PCIe slot on a motherboard of a desktop computer. In at least one embodiment, that PPU may be an integrated graphics processing unit (“iGPU”) included in chipset of a motherboard.

(137) Software compiler logic **150** may be used to perform software compilation operations associated with one or more embodiments, including generating and optimizing peeling loops to execute instructions on serial execution units. Details regarding software compiler logic **150** are provided herein in conjunction with FIG. **1**. In at least one embodiment, software compiler logic **150** and the program code it may generate may be used in the system of FIG. **8** for performing software compilation operations, including generating and optimizing peeling loops to execute instructions on serial execution units, and executing instructions on serial execution units.

(138) FIG. **9** is a block diagram illustrating a computing system **900** according to at least one embodiment. In at least one embodiment, computing system **900** includes a processing subsystem

901 having one or more processor(s) **902** and a system memory **904** communicating via an interconnection path that may include a memory hub **905**. In at least one embodiment, memory hub **905** may be a separate component within a chipset component or may be integrated within one or more processor(s) **902**. In at least one embodiment, memory hub **905** couples with an I/O subsystem **911** via a communication link **906**. In at least one embodiment, I/O subsystem **911** includes an I/O hub **907** that can enable computing system **900** to receive input from one or more input device(s) **908**. In at least one embodiment, I/O hub **907** can enable a display controller, which may be included in one or more processor(s) **902**, to provide outputs to one or more display device(s) **910A**. In at least one embodiment, one or more display device(s) **910A** coupled with I/O hub **907** can include a local, internal, or embedded display device.

(139) In at least one embodiment, processing subsystem **901** includes one or more parallel processor(s) **912** coupled to memory hub **905** via a bus or other communication link **913**. In at least one embodiment, communication link **913** may use one of any number of standards based communication link technologies or protocols, such as, but not limited to PCI Express, or may be a vendor-specific communications interface or communications fabric. In at least one embodiment, one or more parallel processor(s) **912** form a computationally focused parallel or vector processing system that can include a large number of processing cores and/or processing clusters, such as a many-integrated core (MIC) processor. In at least one embodiment, some or all of parallel processor(s) **912** form a graphics processing subsystem that can output pixels to one of one or more display device(s) **910A** coupled via I/O Hub **907**. In at least one embodiment, parallel processor(s) **912** can also include a display controller and display interface (not shown) to enable a direct connection to one or more display device(s) **910B**.

(140) In at least one embodiment, a system storage unit **914** can connect to I/O hub **907** to provide a storage mechanism for computing system **900**. In at least one embodiment, an I/O switch **916** can be used to provide an interface mechanism to enable connections between I/O hub **907** and other components, such as a network adapter **918** and/or a wireless network adapter **919** that may be integrated into platform, and various other devices that can be added via one or more add-in device(s) **920**. In at least one embodiment, network adapter **918** can be an Ethernet adapter or another wired network adapter. In at least one embodiment, wireless network adapter **919** can include one or more of a Wi-Fi, Bluetooth, near field communication (NFC), or other network device that includes one or more wireless radios.

(141) In at least one embodiment, computing system **900** can include other components not explicitly shown, including USB or other port connections, optical storage drives, video capture devices, and like, may also be connected to I/O hub **907**. In at least one embodiment, communication paths interconnecting various components in FIG. 9 may be implemented using any suitable protocols, such as PCI (Peripheral Component Interconnect) based protocols (e.g., PCI-Express), or other bus or point-to-point communication interfaces and/or protocol(s), such as NV-Link high-speed interconnect, or interconnect protocols.

(142) In at least one embodiment, parallel processor(s) **912** incorporate circuitry optimized for graphics and video processing, including, for example, video output circuitry, and constitutes a graphics processing unit (GPU). In at least one embodiment, parallel processor(s) **912** incorporate circuitry optimized for general purpose processing. In at least one embodiment, components of computing system **900** may be integrated with one or more other system elements on a single integrated circuit. For example, in at least one embodiment, parallel processor(s) **912**, memory hub **905**, processor(s) **902**, and I/O hub **907** can be integrated into a system on chip (SoC) integrated circuit. In at least one embodiment, components of computing system **900** can be integrated into a single package to form a system in package (SIP) configuration. In at least one embodiment, at least a portion of components of computing system **900** can be integrated into a multi-chip module (MCM), which can be interconnected with other multi-chip modules into a modular computing system.

(143) Software compiler logic **150** may be used to perform software compilation operations associated with one or more embodiments, including generating and optimizing peeling loops to execute instructions on serial execution units. Details regarding software compiler logic **150** are provided herein in conjunction with FIG. **1**. In at least one embodiment, software compiler logic **150** and the program code it may generate may be used in the system of FIG. **9** for performing software compilation operations, including generating and optimizing peeling loops to execute instructions on serial execution units, and executing instructions on serial execution units.

(144) FIG. **10A** illustrates a parallel processor **1000** according to at least one embodiment. In at least one embodiment, various components of parallel processor **1000** may be implemented using one or more integrated circuit devices, such as programmable processors, application specific integrated circuits (ASICs), or field programmable gate arrays (FPGA). In at least one embodiment, illustrated parallel processor **1000** is a variant of the one or more parallel processor(s) **912** shown in FIG. **9**.

(145) In at least one embodiment, parallel processor **1000** includes a parallel processing unit **1002**. In at least one embodiment, parallel processing unit **1002** includes an I/O unit **1004** that enables communication with other devices, including other instances of parallel processing unit **1002**. In at least one embodiment, I/O unit **1004** may be directly connected to other devices. In at least one embodiment, I/O unit **1004** connects with other devices via use of a hub or switch interface, such as a memory hub **1005**. In at least one embodiment, connections between memory hub **1005** and I/O unit **1004** form a communication link **1013**. In at least one embodiment, I/O unit **1004** connects with a host interface **1006** and a memory crossbar **1016**, where host interface **1006** receives commands directed to performing processing operations and memory crossbar **1016** receives commands directed to performing memory operations.

(146) In at least one embodiment, when host interface **1006** receives a command buffer via I/O unit **1004**, host interface **1006** can direct operations to perform those commands to a front end **1008**. In at least one embodiment, front end **1008** couples with a scheduler **1010**, which is configured to distribute commands or other work items to a processing cluster array **1012**. In at least one embodiment, scheduler **1010** ensures that processing cluster array **1012** is properly configured and in a valid state before tasks are distributed to a cluster of processing cluster array **1012**. In at least one embodiment, scheduler **1010** is implemented via firmware logic executing on a microcontroller. In at least one embodiment, microcontroller implemented scheduler **1010** is configurable to perform complex scheduling and work distribution operations at coarse and fine granularity, e.g., enabling rapid preemption and context switching of threads executing on processing array **1012**. In at least one embodiment, host software can prove workloads for scheduling on processing cluster array **1012** via one of multiple graphics processing paths. In at least one embodiment, workloads can then be automatically distributed across processing array cluster **1012** by scheduler **1010** logic within a microcontroller including scheduler **1010**.

(147) In at least one embodiment, processing cluster array **1012** can include up to “N” processing clusters (e.g., cluster **1014A**, cluster **1014B**, through cluster **1014N**), where “N” represents a positive integer (which may be a different integer “N” than used in other figures). In at least one embodiment, each cluster **1014A-1014N** of processing cluster array **1012** can execute a large number of concurrent threads. In at least one embodiment, scheduler **1010** can allocate work to clusters **1014A-1014N** of processing cluster array **1012** using various scheduling and/or work distribution algorithms, which may vary depending on workload arising for each type of program or computation. In at least one embodiment, scheduling can be handled dynamically by scheduler **1010**, or can be assisted in part by compiler logic during compilation of program logic configured for execution by processing cluster array **1012**. In at least one embodiment, different clusters **1014A-1014N** of processing cluster array **1012** can be allocated for processing different types of programs or for performing different types of computations.

(148) In at least one embodiment, processing cluster array **1012** can be configured to perform

various types of parallel processing operations. In at least one embodiment, processing cluster array **1012** is configured to perform general-purpose parallel compute operations. For example, in at least one embodiment, processing cluster array **1012** can include logic to execute processing tasks including filtering of video and/or audio data, performing modeling operations, including physics operations, and performing data transformations.

(149) In at least one embodiment, processing cluster array **1012** is configured to perform parallel graphics processing operations. In at least one embodiment, processing cluster array **1012** can include additional logic to support execution of such graphics processing operations, including but not limited to, texture sampling logic to perform texture operations, as well as tessellation logic and other vertex processing logic. In at least one embodiment, processing cluster array **1012** can be configured to execute graphics processing related shader programs, for example, such as vertex shaders, tessellation shaders, geometry shaders, and pixel shaders. In at least one embodiment, parallel processing unit **1002** can transfer data from system memory via I/O unit **1004** for processing. In at least one embodiment, during processing, transferred data can be stored to on-chip memory (e.g., parallel processor memory **1022**) during processing, then written back to system memory.

(150) In at least one embodiment, when parallel processing unit **1002** is used to perform graphics processing, scheduler **1010** can be configured to divide a processing workload into approximately equal sized tasks, to better enable distribution of graphics processing operations to multiple clusters **1014A-1014N** of processing cluster array **1012**. In at least one embodiment, portions of processing cluster array **1012** can be configured to perform different types of processing. For example, in at least one embodiment, a first portion may be configured to perform vertex shading and topology generation, a second portion may be configured to perform tessellation and geometry shading, and a third portion may be configured to perform pixel shading or other screen space operations, to produce a rendered image for display. In at least one embodiment, intermediate data produced by one or more of clusters **1014A-1014N** may be stored in buffers to allow intermediate data to be transmitted between clusters **1014A-1014N** for further processing.

(151) In at least one embodiment, processing cluster array **1012** can receive processing tasks to be executed via scheduler **1010**, which receives commands defining processing tasks from front end **1008**. In at least one embodiment, processing tasks can include indices of data to be processed, e.g., surface (patch) data, primitive data, vertex data, and/or pixel data, as well as state parameters and commands defining how data is to be processed (e.g., what program is to be executed). In at least one embodiment, scheduler **1010** may be configured to fetch indices corresponding to tasks or may receive indices from front end **1008**. In at least one embodiment, front end **1008** can be configured to ensure processing cluster array **1012** is configured to a valid state before a workload specified by incoming command buffers (e.g., batch-buffers, push buffers, etc.) is initiated.

(152) In at least one embodiment, each of one or more instances of parallel processing unit **1002** can couple with a parallel processor memory **1022**. In at least one embodiment, parallel processor memory **1022** can be accessed via memory crossbar **1016**, which can receive memory requests from processing cluster array **1012** as well as I/O unit **1004**. In at least one embodiment, memory crossbar **1016** can access parallel processor memory **1022** via a memory interface **1018**. In at least one embodiment, memory interface **1018** can include multiple partition units (e.g., partition unit **1020A**, partition unit **1020B**, through partition unit **1020N**) that can each couple to a portion (e.g., memory unit) of parallel processor memory **1022**. In at least one embodiment, a number of partition units **1020A-1020N** is configured to be equal to a number of memory units, such that a first partition unit **1020A** has a corresponding first memory unit **1024A**, a second partition unit **1020B** has a corresponding memory unit **1024B**, and an N-th partition unit **1020N** has a corresponding N-th memory unit **1024N**. In at least one embodiment, a number of partition units **1020A-1020N** may not be equal to a number of memory units.

(153) In at least one embodiment, memory units **1024A-1024N** can include various types of

memory devices, including dynamic random access memory (DRAM) or graphics random access memory, such as synchronous graphics random access memory (SGRAM), including graphics double data rate (GDDR) memory. In at least one embodiment, memory units **1024A-1024N** may also include 3D stacked memory, including but not limited to high bandwidth memory (HBM). In at least one embodiment, render targets, such as frame buffers or texture maps may be stored across memory units **1024A-1024N**, allowing partition units **1020A-1020N** to write portions of each render target in parallel to efficiently use available bandwidth of parallel processor memory **1022**. In at least one embodiment, a local instance of parallel processor memory **1022** may be excluded in favor of a unified memory design that utilizes system memory in conjunction with local cache memory.

(154) In at least one embodiment, any one of clusters **1014A-1014N** of processing cluster array **1012** can process data that will be written to any of memory units **1024A-1024N** within parallel processor memory **1022**. In at least one embodiment, memory crossbar **1016** can be configured to transfer an output of each cluster **1014A-1014N** to any partition unit **1020A-1020N** or to another cluster **1014A-1014N**, which can perform additional processing operations on an output. In at least one embodiment, each cluster **1014A-1014N** can communicate with memory interface **1018** through memory crossbar **1016** to read from or write to various external memory devices. In at least one embodiment, memory crossbar **1016** has a connection to memory interface **1018** to communicate with I/O unit **1004**, as well as a connection to a local instance of parallel processor memory **1022**, enabling processing units within different processing clusters **1014A-1014N** to communicate with system memory or other memory that is not local to parallel processing unit **1002**. In at least one embodiment, memory crossbar **1016** can use virtual channels to separate traffic streams between clusters **1014A-1014N** and partition units **1020A-1020N**.

(155) In at least one embodiment, multiple instances of parallel processing unit **1002** can be provided on a single add-in card, or multiple add-in cards can be interconnected. In at least one embodiment, different instances of parallel processing unit **1002** can be configured to interoperate even if different instances have different numbers of processing cores, different amounts of local parallel processor memory, and/or other configuration differences. For example, in at least one embodiment, some instances of parallel processing unit **1002** can include higher precision floating point units relative to other instances. In at least one embodiment, systems incorporating one or more instances of parallel processing unit **1002** or parallel processor **1000** can be implemented in a variety of configurations and form factors, including but not limited to desktop, laptop, or handheld personal computers, servers, workstations, game consoles, and/or embedded systems.

(156) FIG. **10B** is a block diagram of a partition unit **1020** according to at least one embodiment. In at least one embodiment, partition unit **1020** is an instance of one of partition units **1020A-1020N** of FIG. **10A**. In at least one embodiment, partition unit **1020** includes an L2 cache **1021**, a frame buffer interface **1025**, and a ROP **1026** (raster operations unit). In at least one embodiment, L2 cache **1021** is a read/write cache that is configured to perform load and store operations received from memory crossbar **1016** and ROP **1026**. In at least one embodiment, read misses and urgent write-back requests are output by L2 cache **1021** to frame buffer interface **1025** for processing. In at least one embodiment, updates can also be sent to a frame buffer via frame buffer interface **1025** for processing. In at least one embodiment, frame buffer interface **1025** interfaces with one of memory units in parallel processor memory, such as memory units **1024A-1024N** of FIG. **10** (e.g., within parallel processor memory **1022**).

(157) In at least one embodiment, ROP **1026** is a processing unit that performs raster operations such as stencil, z test, blending, etc. In at least one embodiment, ROP **1026** then outputs processed graphics data that is stored in graphics memory. In at least one embodiment, ROP **1026** includes compression logic to compress depth or color data that is written to memory and decompress depth or color data that is read from memory. In at least one embodiment, compression logic can be lossless compression logic that makes use of one or more of multiple compression algorithms. In at

least one embodiment, a type of compression that is performed by ROP **1026** can vary based on statistical characteristics of data to be compressed. For example, in at least one embodiment, delta color compression is performed on depth and color data on a per-tile basis.

(158) In at least one embodiment, ROP **1026** is included within each processing cluster (e.g., cluster **1014A-1014N** of FIG. **10A**) instead of within partition unit **1020**. In at least one embodiment, read and write requests for pixel data are transmitted over memory crossbar **1016** instead of pixel fragment data. In at least one embodiment, processed graphics data may be displayed on a display device, such as one of one or more display device(s) **910** of FIG. **9**, routed for further processing by processor(s) **902**, or routed for further processing by one of processing entities within parallel processor **1000** of FIG. **10A**.

(159) FIG. **10C** is a block diagram of a processing cluster **1014** within a parallel processing unit according to at least one embodiment. In at least one embodiment, a processing cluster is an instance of one of processing clusters **1014A-1014N** of FIG. **10A**. In at least one embodiment, processing cluster **1014** can be configured to execute many threads in parallel, where “thread” refers to an instance of a particular program executing on a particular set of input data. In at least one embodiment, single-instruction, multiple-data (SIMD) instruction issue techniques are used to support parallel execution of a large number of threads without providing multiple independent instruction units. In at least one embodiment, single-instruction, multiple-thread (SIMT) techniques are used to support parallel execution of a large number of generally synchronized threads, using a common instruction unit configured to issue instructions to a set of processing engines within each one of processing clusters.

(160) In at least one embodiment, operation of processing cluster **1014** can be controlled via a pipeline manager **1032** that distributes processing tasks to SIMT parallel processors. In at least one embodiment, pipeline manager **1032** receives instructions from scheduler **1010** of FIG. **10A** and manages execution of those instructions via a graphics multiprocessor **1034** and/or a texture unit **1036**. In at least one embodiment, graphics multiprocessor **1034** is an exemplary instance of a SIMT parallel processor. However, in at least one embodiment, various types of SIMT parallel processors of differing architectures may be included within processing cluster **1014**. In at least one embodiment, one or more instances of graphics multiprocessor **1034** can be included within a processing cluster **1014**. In at least one embodiment, graphics multiprocessor **1034** can process data and a data crossbar **1040** can be used to distribute processed data to one of multiple possible destinations, including other shader units. In at least one embodiment, pipeline manager **1032** can facilitate distribution of processed data by specifying destinations for processed data to be distributed via data crossbar **1040**.

(161) In at least one embodiment, each graphics multiprocessor **1034** within processing cluster **1014** can include an identical set of functional execution logic (e.g., arithmetic logic units, load-store units, etc.). In at least one embodiment, functional execution logic can be configured in a pipelined manner in which new instructions can be issued before previous instructions are complete. In at least one embodiment, functional execution logic supports a variety of operations including integer and floating point arithmetic, comparison operations, Boolean operations, bit-shifting, and computation of various algebraic functions. In at least one embodiment, same functional-unit hardware can be leveraged to perform different operations and any combination of functional units may be present.

(162) In at least one embodiment, instructions transmitted to processing cluster **1014** constitute a thread. In at least one embodiment, a set of threads executing across a set of parallel processing engines is a thread group. In at least one embodiment, a thread group executes a common program on different input data. In at least one embodiment, each thread within a thread group can be assigned to a different processing engine within a graphics multiprocessor **1034**. In at least one embodiment, a thread group may include fewer threads than a number of processing engines within graphics multiprocessor **1034**. In at least one embodiment, when a thread group includes fewer

threads than a number of processing engines, one or more of processing engines may be idle during cycles in which that thread group is being processed. In at least one embodiment, a thread group may also include more threads than a number of processing engines within graphics multiprocessor **1034**. In at least one embodiment, when a thread group includes more threads than number of processing engines within graphics multiprocessor **1034**, processing can be performed over consecutive clock cycles. In at least one embodiment, multiple thread groups can be executed concurrently on a graphics multiprocessor **1034**.

(163) In at least one embodiment, graphics multiprocessor **1034** includes an internal cache memory to perform load and store operations. In at least one embodiment, graphics multiprocessor **1034** can forego an internal cache and use a cache memory (e.g., L1 cache **1048**) within processing cluster **1014**. In at least one embodiment, each graphics multiprocessor **1034** also has access to L2 caches within partition units (e.g., partition units **1020A-1020N** of FIG. **10A**) that are shared among all processing clusters **1014** and may be used to transfer data between threads. In at least one embodiment, graphics multiprocessor **1034** may also access off-chip global memory, which can include one or more of local parallel processor memory and/or system memory. In at least one embodiment, any memory external to parallel processing unit **1002** may be used as global memory. In at least one embodiment, processing cluster **1014** includes multiple instances of graphics multiprocessor **1034** and can share common instructions and data, which may be stored in L1 cache **1048**.

(164) In at least one embodiment, each processing cluster **1014** may include an MMU **1045** (memory management unit) that is configured to map virtual addresses into physical addresses. In at least one embodiment, one or more instances of MMU **1045** may reside within memory interface **1018** of FIG. **10A**. In at least one embodiment, MMU **1045** includes a set of page table entries (PTEs) used to map a virtual address to a physical address of a tile and optionally a cache line index. In at least one embodiment, MMU **1045** may include address translation lookaside buffers (TLB) or caches that may reside within graphics multiprocessor **1034** or L1 **1048** cache or processing cluster **1014**. In at least one embodiment, a physical address is processed to distribute surface data access locally to allow for efficient request interleaving among partition units. In at least one embodiment, a cache line index may be used to determine whether a request for a cache line is a hit or miss.

(165) In at least one embodiment, a processing cluster **1014** may be configured such that each graphics multiprocessor **1034** is coupled to a texture unit **1036** for performing texture mapping operations, e.g., determining texture sample positions, reading texture data, and filtering texture data. In at least one embodiment, texture data is read from an internal texture L1 cache (not shown) or from an L1 cache within graphics multiprocessor **1034** and is fetched from an L2 cache, local parallel processor memory, or system memory, as needed. In at least one embodiment, each graphics multiprocessor **1034** outputs processed tasks to data crossbar **1040** to provide processed task to another processing cluster **1014** for further processing or to store processed task in an L2 cache, local parallel processor memory, or system memory via memory crossbar **1016**. In at least one embodiment, a preROP **1042** (pre-raster operations unit) is configured to receive data from graphics multiprocessor **1034**, and direct data to ROP units, which may be located with partition units as described herein (e.g., partition units **1020A-1020N** of FIG. **10A**). In at least one embodiment, preROP **1042** unit can perform optimizations for color blending, organizing pixel color data, and performing address translations.

(166) Software compiler logic **150** may be used to perform software compilation operations associated with one or more embodiments, including generating and optimizing peeling loops to execute instructions on serial execution units. Details regarding software compiler logic **150** are provided herein in conjunction with FIG. **1**. In at least one embodiment, software compiler logic **150** and the program code it may generate may be used in the system of FIG. **10C** for performing software compilation operations, including generating and optimizing peeling loops to execute

instructions on serial execution units, and executing instructions on serial execution units.

(167) FIG. 10D shows a graphics multiprocessor **1034** according to at least one embodiment. In at least one embodiment, graphics multiprocessor **1034** couples with pipeline manager **1032** of processing cluster **1014**. In at least one embodiment, graphics multiprocessor **1034** has an execution pipeline including but not limited to an instruction cache **1052**, an instruction unit **1054**, an address mapping unit **1056**, a register **1058**, one or more general purpose graphics processing unit (GPGPU) cores **1062**, and one or more load/store units **1066**. In at least one embodiment, GPGPU cores **1062** and load/store units **1066** are coupled with cache memory **1072** and shared memory **1070** via a memory and cache interconnect **1068**.

(168) In at least one embodiment, instruction cache **1052** receives a stream of instructions to execute from pipeline manager **1032**. In at least one embodiment, instructions are cached in instruction cache **1052** and dispatched for execution by an instruction unit **1054**. In at least one embodiment, instruction unit **1054** can dispatch instructions as thread groups (e.g., warps), with each thread of thread group assigned to a different execution unit within GPGPU cores **1062**. In at least one embodiment, an instruction can access any of a local, shared, or global address space by specifying an address within a unified address space. In at least one embodiment, address mapping unit **1056** can be used to translate addresses in a unified address space into a distinct memory address that can be accessed by load/store units **1066**.

(169) In at least one embodiment, register **1058** provides a set of registers for functional units of graphics multiprocessor **1034**. In at least one embodiment, register **1058** provides temporary storage for operands connected to data paths of functional units (e.g., GPGPU cores **1062**, load/store units **1066**) of graphics multiprocessor **1034**. In at least one embodiment, register **1058** is divided between each of functional units such that each functional unit is allocated a dedicated portion of register **1058**. In at least one embodiment, register **1058** is divided between different warps being executed by graphics multiprocessor **1034**.

(170) In at least one embodiment, GPGPU cores **1062** can each include floating point units (FPUs) and/or integer arithmetic logic units (ALUs) that are used to execute instructions of graphics multiprocessor **1034**. In at least one embodiment, GPGPU cores **1062** can be similar in architecture or can differ in architecture. In at least one embodiment, a first portion of GPGPU cores **1062** include a single precision FPU and an integer ALU while a second portion of GPGPU cores include a double precision FPU. In at least one embodiment, FPUs can implement IEEE 754-2008 standard floating point arithmetic or enable variable precision floating point arithmetic. In at least one embodiment, graphics multiprocessor **1034** can additionally include one or more fixed function or special function units to perform specific functions such as copy rectangle or pixel blending operations. In at least one embodiment, one or more of GPGPU cores **1062** can also include fixed or special function logic.

(171) In at least one embodiment, GPGPU cores **1062** include SIMD logic capable of performing a single instruction on multiple sets of data. In at least one embodiment, GPGPU cores **1062** can physically execute SIMD4, SIMD8, and SIMD16 instructions and logically execute SIMD1, SIMD2, and SIMD32 instructions. In at least one embodiment, SIMD instructions for GPGPU cores can be generated at compile time by a shader compiler or automatically generated when executing programs written and compiled for single program multiple data (SPMD) or SIMT architectures. In at least one embodiment, multiple threads of a program configured for an SIMT execution model can be executed via a single SIMD instruction. For example, in at least one embodiment, eight SIMT threads that perform same or similar operations can be executed in parallel via a single SIMD8 logic unit.

(172) In at least one embodiment, memory and cache interconnect **1068** is an interconnect network that connects each functional unit of graphics multiprocessor **1034** to register **1058** and to shared memory **1070**. In at least one embodiment, memory and cache interconnect **1068** is a crossbar interconnect that allows load/store unit **1066** to implement load and store operations between

shared memory **1070** and register **1058**. In at least one embodiment, register **1058** can operate at a same frequency as GPGPU cores **1062**, thus data transfer between GPGPU cores **1062** and register **1058** can have very low latency. In at least one embodiment, shared memory **1070** can be used to enable communication between threads that execute on functional units within graphics multiprocessor **1034**. In at least one embodiment, cache memory **1072** can be used as a data cache for example, to cache texture data communicated between functional units and texture unit **1036**. In at least one embodiment, shared memory **1070** can also be used as a program managed cache. In at least one embodiment, threads executing on GPGPU cores **1062** can programmatically store data within shared memory in addition to automatically cached data that is stored within cache memory **1072**.

(173) In at least one embodiment, a parallel processor or GPGPU as described herein is communicatively coupled to host/processor cores to accelerate graphics operations, machine-learning operations, pattern analysis operations, and various general purpose GPU (GPGPU) functions. In at least one embodiment, a GPU may be communicatively coupled to host processor/cores over a bus or other interconnect (e.g., a high-speed interconnect such as PCIe or NVLink). In at least one embodiment, a GPU may be integrated on a package or chip as cores and communicatively coupled to cores over an internal processor bus/interconnect internal to a package or chip. In at least one embodiment, regardless a manner in which a GPU is connected, processor cores may allocate work to such GPU in a form of sequences of commands/instructions contained in a work descriptor. In at least one embodiment, that GPU then uses dedicated circuitry/logic for efficiently processing these commands/instructions.

(174) Software compiler logic **150** may be used to perform software compilation operations associated with one or more embodiments, including generating and optimizing peeling loops to execute instructions on serial execution units. Details regarding software compiler logic **150** are provided herein in conjunction with FIG. **1**. In at least one embodiment, software compiler logic **150** and the program code it may generate may be used in the system of FIG. **10D** for performing software compilation operations, including generating and optimizing peeling loops to execute instructions on serial execution units, and executing instructions on serial execution units.

(175) Other variations are within spirit of present disclosure. Thus, while disclosed techniques are susceptible to various modifications and alternative constructions, certain illustrated embodiments thereof are shown in drawings and have been described above in detail. It should be understood, however, that there is no intention to limit disclosure to specific form or forms disclosed, but on contrary, intention is to cover all modifications, alternative constructions, and equivalents falling within spirit and scope of disclosure, as defined in appended claims.

(176) Use of terms “a” and “an” and “the” and similar referents in context of describing disclosed embodiments (especially in context of following claims) are to be construed to cover both singular and plural, unless otherwise indicated herein or clearly contradicted by context, and not as a definition of a term. Terms “comprising,” “having,” “including,” and “containing” are to be construed as open-ended terms (meaning “including, but not limited to,”) unless otherwise noted. “Connected,” when unmodified and referring to physical connections, is to be construed as partly or wholly contained within, attached to, or joined together, even if there is something intervening. Recitation of ranges of values herein are merely intended to serve as a shorthand method of referring individually to each separate value falling within range, unless otherwise indicated herein and each separate value is incorporated into specification as if it were individually recited herein. In at least one embodiment, use of term “set” (e.g., “a set of items”) or “subset” unless otherwise noted or contradicted by context, is to be construed as a nonempty collection comprising one or more members. Further, unless otherwise noted or contradicted by context, term “subset” of a corresponding set does not necessarily denote a proper subset of corresponding set, but subset and corresponding set may be equal.

(177) Conjunctive language, such as phrases of form “at least one of A, B, and C,” or “at least one

of A, B and C,” unless specifically stated otherwise or otherwise clearly contradicted by context, is otherwise understood with context as used in general to present that an item, term, etc., may be either A or B or C, or any nonempty subset of set of A and B and C. For instance, in illustrative example of a set having three members, conjunctive phrases “at least one of A, B, and C” and “at least one of A, B and C” refer to any of following sets: {A}, {B}, {C}, {A, B}, {A, C}, {B, C}, {A, B, C}. Thus, such conjunctive language is not generally intended to imply that certain embodiments require at least one of A, at least one of B and at least one of C each to be present. In addition, unless otherwise noted or contradicted by context, term “plurality” indicates a state of being plural (e.g., “a plurality of items” indicates multiple items). In at least one embodiment, number of items in a plurality is at least two, but can be more when so indicated either explicitly or by context. Further, unless stated otherwise or otherwise clear from context, phrase “based on” means “based at least in part on” and not “based solely on.”

(178) Operations of processes described herein can be performed in any suitable order unless otherwise indicated herein or otherwise clearly contradicted by context. In at least one embodiment, a process such as those processes described herein (or variations and/or combinations thereof) is performed under control of one or more computer systems configured with executable instructions and is implemented as code (e.g., executable instructions, one or more computer programs or one or more applications) executing collectively on one or more processors, by hardware or combinations thereof. In at least one embodiment, code is stored on a computer-readable storage medium, for example, in form of a computer program comprising a plurality of instructions executable by one or more processors. In at least one embodiment, a computer-readable storage medium is a non-transitory computer-readable storage medium that excludes transitory signals (e.g., a propagating transient electric or electromagnetic transmission) but includes non-transitory data storage circuitry (e.g., buffers, cache, and queues) within transceivers of transitory signals. In at least one embodiment, code (e.g., executable code or source code) is stored on a set of one or more non-transitory computer-readable storage media having stored thereon executable instructions (or other memory to store executable instructions) that, when executed (e.g., as a result of being executed) by one or more processors of a computer system, cause computer system to perform operations described herein. In at least one embodiment, set of non-transitory computer-readable storage media comprises multiple non-transitory computer-readable storage media and one or more of individual non-transitory storage media of multiple non-transitory computer-readable storage media lack all of code while multiple non-transitory computer-readable storage media collectively store all of code. In at least one embodiment, executable instructions are executed such that different instructions are executed by different processors for example, a non-transitory computer-readable storage medium store instructions and a main central processing unit (“CPU”) executes some of instructions while a graphics processing unit (“GPU”) executes other instructions. In at least one embodiment, different components of a computer system have separate processors and different processors execute different subsets of instructions.

(179) Accordingly, in at least one embodiment, computer systems are configured to implement one or more services that singly or collectively perform operations of processes described herein and such computer systems are configured with applicable hardware and/or software that enable performance of operations. Further, a computer system that implements at least one embodiment of present disclosure is a single device and, in another embodiment, is a distributed computer system comprising multiple devices that operate differently such that distributed computer system performs operations described herein and such that a single device does not perform all operations.

(180) Use of any and all examples, or exemplary language (e.g., “such as”) provided herein, is intended merely to better illuminate embodiments of disclosure and does not pose a limitation on scope of disclosure unless otherwise claimed. No language in specification should be construed as indicating any non-claimed element as essential to practice of disclosure.

(181) All references, including publications, patent applications, and patents, cited herein are hereby incorporated by reference to same extent as if each reference were individually and specifically indicated to be incorporated by reference and were set forth in its entirety herein.

(182) In description and claims, terms “coupled” and “connected,” along with their derivatives, may be used. It should be understood that these terms may be not intended as synonyms for each other. Rather, in particular examples, “connected” or “coupled” may be used to indicate that two or more elements are in direct or indirect physical or electrical contact with each other. “Coupled” may also mean that two or more elements are not in direct contact with each other, but yet still co-operate or interact with each other.

(183) Unless specifically stated otherwise, it may be appreciated that throughout specification terms such as “processing,” “computing,” “calculating,” “determining,” or like, refer to action and/or processes of a computer or computing system, or similar electronic computing device, that manipulate and/or transform data represented as physical, such as electronic, quantities within computing system's registers and/or memories into other data similarly represented as physical quantities within computing system's memories, registers or other such information storage, transmission or display devices.

(184) In a similar manner, term “processor” may refer to any device or portion of a device that processes electronic data from registers and/or memory and transform that electronic data into other electronic data that may be stored in registers and/or memory. As non-limiting examples, “processor” may be a CPU or a GPU. A “computing platform” may comprise one or more processors. As used herein, “software” processes may include, for example, software and/or hardware entities that perform work over time, such as tasks, threads, and intelligent agents. Also, each process may refer to multiple processes, for carrying out instructions in sequence or in parallel, continuously or intermittently. In at least one embodiment, terms “system” and “method” are used herein interchangeably insofar as system may embody one or more methods and methods may be considered a system.

(185) In present document, references may be made to obtaining, acquiring, receiving, or inputting analog or digital data into a subsystem, computer system, or computer-implemented machine. In at least one embodiment, process of obtaining, acquiring, receiving, or inputting analog and digital data can be accomplished in a variety of ways such as by receiving data as a parameter of a function call or a call to an application programming interface. In at least one embodiment, processes of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a serial or parallel interface. In at least one embodiment, processes of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a computer network from providing entity to acquiring entity. In at least one embodiment, references may also be made to providing, outputting, transmitting, sending, or presenting analog or digital data. In various examples, processes of providing, outputting, transmitting, sending, or presenting analog or digital data can be accomplished by transferring data as an input or output parameter of a function call, a parameter of an application programming interface or interprocess communication mechanism.

(186) Although descriptions herein set forth example implementations of described techniques, other architectures may be used to implement described functionality, and are intended to be within scope of this disclosure. Furthermore, although specific distributions of responsibilities may be defined above for purposes of description, various functions and responsibilities might be distributed and divided in different ways, depending on circumstances.

(187) Furthermore, although subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that subject matter claimed in appended claims is not necessarily limited to specific features or acts described. Rather, specific features and acts are disclosed as exemplary forms of implementing the claims.

Claims

1. A processor comprising: a parallel execution unit; and a serial execution unit, wherein the processor operates to: select a set of one or more threads from a group of active threads associated with an instruction; copy a source operand value associated with a thread of the set of one or more threads from a first private register to a first shared register to obtain a shared source operand; and execute the instruction for the set of one or more threads using the serial execution unit and based at least on the shared source operand.
2. The processor of claim 1, wherein the set of one or more threads are selected based at least on the set of one or more threads having a common source operand value.
3. The processor of claim 2, wherein the processor further operates to: select, from the set of one or more threads, a subset of one or more threads having an additional common source operand value; and execute the instruction for the subset of one or more threads using the serial execution unit.
4. The processor of claim 2, wherein the processor further operates to: copy the common source operand value from a second private register to a second shared register to obtain a second shared source operand; execute the instruction for the set of one or more threads using the serial execution unit and based at least on the second shared source operand; and store a result in a destination shared register.
5. The processor of claim 4, wherein the processor further operates to: select, from a remaining group of active threads associated with the instruction that are not included in the set of one or more threads, another set of one or more threads having another common source operand value; copy the another common source operand value from another private register to another shared register to obtain a third shared source operand; execute the instruction on the serial execution unit based at least on the third shared source operand; and store another result in another destination shared register.
6. The processor of claim 4, wherein the processor further operates to: copy the result to a destination private register for at least one individual thread of the set of one or more threads.
7. The processor of claim 1, wherein the set of one or more threads includes a single thread, and wherein the processor further operates to: store a result in a destination shared register.
8. The processor of claim 7, wherein the processor further operates to: select another thread from a remaining group of active threads associated with the instruction; copy another source operand value of the another thread from another private register to another shared register to obtain another shared source operand; and execute the instruction for the another thread using the serial execution unit and based at least on the another shared source operand.
9. The processor of claim 1, wherein the processor further operates to: execute a different instruction for at least one active thread of the group of active threads using the parallel execution unit.
10. The processor of claim 1, wherein the processor is comprised in at least one of: a control system for an autonomous or semi-autonomous machine; a perception system for an autonomous or semi-autonomous machine; a system for performing simulation operations; a system for performing digital twin operations; a system for performing light transport simulation; a system for performing collaborative content creation for 3D assets; a system for performing deep learning operations; a system for generating or presenting at least one of augmented reality content, virtual reality content, or mixed reality content; a system for hosting one or more real-time streaming applications; a system implemented using an edge device; a system implemented using a robot; a system for performing conversational AI operations; a system for generating synthetic data; a system incorporating one or more virtual machines (VMs); a system implemented at least partially in a data center; or a system implemented at least partially using cloud computing resources.
11. A system comprising: one or more processing units to: replace an instruction associated with a

group of threads and corresponding to a parallel execution unit with one or more instructions to: select, from the group of threads, a set of one or more threads having a common source operand value; copy the common source operand value associated with a thread of the set of one or more threads from a private register to a shared register to obtain a shared source operand; and execute the instruction for the set of one or more threads using a serial execution unit and based at least on the shared source operand.

12. The system of claim 11, wherein the one or more instructions are further to: store a result in a shared destination register.

13. The system of claim 11, wherein the one or more instructions are further to: select another thread from a remaining group of active threads not included in the set of one or more threads; copy another source operand value of the another thread from another private register to another shared register to obtain another shared source operand; and execute the instruction for the another thread on the serial execution unit using the another shared source operand.

14. The system of claim 11, wherein the group of threads is identified using a predicate mask, and wherein the predicate mask is a predicate guard or a source operand for at least one instruction of the one or more instructions.

15. The system of claim 14, wherein the one or more instructions are further to: initialize the predicate mask with an inherited predicate of the instruction.

16. The system of claim 14, wherein the one or more instructions are further to: update the predicate mask to indicate that the instruction is no longer pending execution for the set of one or more threads.

17. The system of claim 11, wherein the one or more processing units are further to: analyze a program code including the instruction to determine a likelihood of data uniformity with respect to a set of source operands of the instruction; and responsive to a determination that the likelihood satisfies a threshold criteria, replace the instruction with the one or more instructions.

18. The system of claim 11, wherein the system is comprised in at least one of: a control system for an autonomous or semi-autonomous machine; a perception system for an autonomous or semi-autonomous machine; a system for performing simulation operations; a system for performing digital twin operations; a system for performing light transport simulation; a system for performing collaborative content creation for 3D assets; a system for performing deep learning operations; a system for generating or presenting at least one of augmented reality content, virtual reality content, or mixed reality content; a system for hosting one or more real-time streaming applications; a system implemented using an edge device; a system implemented using a robot; a system for performing conversational AI operations; a system for generating synthetic data; a system incorporating one or more virtual machines (VMs); a system implemented at least partially in a data center; or a system implemented at least partially using cloud computing resources.

19. A method comprising: copying a same operand value associated with a set of one or more threads from a private register to a shared register to obtain a shared source operand; and executing, using a serial execution unit, an instruction associated with the set of one or more threads based at least on the same operand value being associated with the set of one or more threads.

20. The method of claim 19, further comprising: setting one or more predicate flags corresponding to the set of one or more threads; and causing, based at least on the one or more predicate flags, the set of one or more threads to refrain from executing the instruction with respect to the same operand value.
