

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12405802
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Iuliano; Luca et al.

Methods and systems for inter-pipeline data hazard avoidance

Abstract

Methods and parallel processing units for avoiding inter-pipeline data hazards identified at compile time. For each identified inter-pipeline data hazard the primary instruction and secondary instruction(s) thereof are identified as such and are linked by a counter which is used to track that inter-pipeline data hazard. When a primary instruction is output by the instruction decoder for execution the value of the counter associated therewith is adjusted to indicate that there is hazard related to the primary instruction, and when primary instruction has been resolved by one of multiple parallel processing pipelines the value of the counter associated therewith is adjusted to indicate that the hazard related to the primary instruction has been resolved. When a secondary instruction is output by the decoder for execution, the secondary instruction is stalled in a queue associated with the appropriate instruction pipeline if at least one counter associated with the primary instructions from which it depends indicates that there is a hazard related to the primary instruction.

Inventors: Iuliano; Luca (Chesham, GB), Nield; Simon (Wendover, GB), Foo; Yoong-Chert (London, GB), Mower; Ollie (Harrow, GB)

Applicant: Imagination Technologies Limited (Kings Langley, GB)

Family ID: 59462476

Assignee: Imagination Technologies Limited (Kings Langley, GB)

Appl. No.: 18/439700

Filed: February 12, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20250021340 A1	Jan. 16, 2025

Related U.S. Application Data

Publication Classification**Int. Cl.:** G06F9/22 (20060101); G06F9/30 (20180101); G06F9/38 (20180101)**U.S. Cl.:****CPC** G06F9/3861 (20130101); G06F9/3016 (20130101); G06F9/3834 (20130101); G06F9/3838 (20130101); G06F9/3867 (20130101); G06F9/3889 (20130101);**Field of Classification Search****CPC:** G06F (9/3861); G06F (9/3016); G06F (9/3834); G06F (9/3838); G06F (9/3867); G06F (9/3889)

References Cited**U.S. PATENT DOCUMENTS**

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5961630	12/1998	Zaidi et al.	N/A	N/A
6065105	12/1999	Zaidi et al.	N/A	N/A
6112019	12/1999	Chamdani et al.	N/A	N/A
6604190	12/2002	Tran	712/240	G06F 9/3826
6711668	12/2003	Wojcieszak	712/E9.055	G06F 9/30181
6715060	12/2003	Arnold	712/217	G06F 9/3858
6950927	12/2004	Apisdorf et al.	N/A	N/A
7340590	12/2007	Sugumar et al.	N/A	N/A
8214624	12/2011	Berglas et al.	N/A	N/A
8407454	12/2012	Berglas	712/216	G06F 9/3851
8553489	12/2012	Fujisawa	365/230.09	G11C 8/18
9389869	12/2015	Tran	N/A	G06F 9/3858
10176546	12/2018	Nystad	N/A	G06F 9/3888
10467010	12/2018	Abdallah	N/A	G06F 9/30043
11900122	12/2023	Iuliano	N/A	G06F 9/3861
2005/0076189	12/2004	Wittenburg	712/217	G06F 9/3836
2007/0204135	12/2006	Jiang	N/A	N/A
2007/0239975	12/2006	Wang	712/241	G06F 9/3861
2007/0260856	12/2006	Tran	712/217	G06F 9/3836
2011/0055531	12/2010	Bellows et al.	N/A	N/A
2011/0153986	12/2010	Alexander et al.	N/A	N/A
2022/0066781	12/2021	Iuliano	N/A	G06F 9/3838

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
103377037	12/2012	CN	N/A
2447907	12/2007	GB	N/A
2514618	12/2013	GB	N/A
2554392	12/2018	GB	N/A

6-89205	12/1993	JP	N/A
87/02799	12/1986	WO	N/A
93/20505	12/1992	WO	N/A

OTHER PUBLICATIONS

Lee et al., “Dynamically Scheduling VLIW Instructions with Dependency Information,” Proceedings of the Sixth Annual Workshop on Interaction between Compilers and Computer Architectures, IEEE 2002, pp. 15-23. cited by applicant

“Register Dataflow,” Superscalar Microprocessor Design, Jan. 1, 1991, pp. 103-126. cited by applicant

*(Note: NPL in parent app'n). cited by applicant

Primary Examiner: Tseng; Cheng Yuan

Attorney, Agent or Firm: Potomac Law Group, PLLC

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS AND CLAIM OF PRIORITY (1) This application is a continuation under 35 U.S.C. 120 of copending application Ser. No. 18/220,048 filed Jul. 10, 2023, now U.S. Pat. No. 11,900,122, which is a continuation of prior application Ser. No. 17/523,633 filed Nov. 10, 2021, now U.S. Pat. No. 11,698,790, which is a continuation of prior application Ser. No. 17/070,316 filed Oct. 14, 2020, now U.S. Pat. No. 11,200,064, which is a continuation of prior application Ser. No. 16/009,358 filed Jun. 15, 2018, now U.S. Pat. No. 10,817,301, which claims foreign priority under 35 U.S.C. 119 from United Kingdom Application Nos. 1709598.5 filed Jun. 16, 2017, and 1720408.2 filed Dec. 7, 2017, the contents of which are incorporated herein by reference in their entirety.

BACKGROUND OF THE INVENTION

(1) As is known to those of skill in the art, a data hazard is created in a processing unit with an instruction pipeline when the pipelining of instructions changes the order of read and write accesses to instruction operands so that the order differs from the order that would occur from sequentially executing the instructions one-by-one.

(2) There are three classes of data hazards: read after write (RAW); write after read (WAR); and write after write (WAW)—which are named after the ordering in the program that must be preserved by the pipeline. A RAW data hazard is the most common type of data hazard and occurs when a later instruction (with respect to the order of the instructions in the program) tries to read a source operand before an earlier instruction writes to that source operand. This results in the later instruction getting the old value of the operand. For example, if there is the following set of instructions:

(3) $R1 = R2 + R3$
 $R4 = R1 - R5$

wherein the first instruction causes the sum of the values of register 2 (R2) and register 3 (R3) to be stored in register 1 (R1) and the second instruction causes the difference between the value of register 1 (R1) and register 5 (R5) to be stored in register 4 (R4), a RAW data hazard occurs if the second instruction reads register 1 (R1) before the first instruction has written to register 1 (R1). A WAW data hazard occurs when a later instruction (with respect to the order of the instructions in the program) writes to an operand before it is written to by an earlier instruction which results in the writes being performed in the wrong order so that the operand has the value from the earlier instruction instead of the value from the later instruction. A WAR data hazard occurs when a later

instruction (with respect to the order of the instructions in the program) tries to write to an operand before it is read by an earlier instruction which results in the earlier instruction reading the incorrect value.

(4) There are many known methods, such as forwarding, for avoiding data hazards caused by a single instruction pipeline, however many processing units, such as graphics processing units (GPUs), are configured with a plurality of parallel instruction pipelines to efficiently process large amounts of data in parallel. In such parallel processing units not only do intra-pipeline hazards (i.e. hazards related to instructions that are executed in the same instruction pipeline) need to be tracked and eliminated, but inter-pipeline hazards (i.e. hazards related to instructions that are executed in different instruction pipelines) also need to be tracked and eliminated.

(5) The embodiments described below are provided by way of example only and are not limiting of implementations which solve any or all of the disadvantages of known GPUs or parallel processing units.

SUMMARY OF THE INVENTION

(6) This summary is provided to introduce a selection of concepts that are further described below in the detailed description. This summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used to limit the scope of the claimed subject matter.

(7) Described herein are methods and parallel processing units for avoiding inter-pipeline data hazards where inter-pipeline data hazards are identified at compile time. For each identified inter-pipeline data hazard the primary instruction and secondary instruction(s) thereof are identified and linked by a counter used to track that inter-pipeline data hazard. When a primary instruction is output by the decoder for execution the value of the counter associated therewith is adjusted (e.g. incremented) to indicate a hazard related to that primary instruction, and when it is detected that the hazard related to that primary instruction has been resolved (e.g. the primary instruction has written data to memory) the value of the counter associated therewith is adjusted (e.g. decremented) to indicate that the hazard has been resolved. When a secondary instruction is output by the decoder for execution, the secondary instruction is stalled in a queue associated with the appropriate instruction pipeline if at least one counter associated with a primary instruction from which it depends indicates that there is a hazard related to the primary instruction.

(8) A first aspect provides a parallel processing unit comprising: a plurality of counters; a plurality of queues, each queue preceding one instruction pipeline of a plurality of instruction pipelines; an instruction decoder configured to: decode a received instruction; in response to determining the decoded instruction is a primary instruction from which at least one other instruction is dependent on, cause a value of a counter of the plurality of counters associated with the primary instruction to be adjusted to indicate that there is a hazard related to the primary instruction; and forward the decoded instruction to one of the plurality of queues; and monitor logic configured to monitor the plurality of instruction pipelines, and in response to detecting that an instruction pipeline has resolved a hazard related to a primary instruction, cause the value of the counter associated with the primary instruction to be adjusted to indicate that the hazard related to the primary instruction has been resolved; wherein each queue is configured to, in response to receiving a secondary instruction that is dependent on one or more primary instructions, stall execution of the secondary instruction by the associated instruction pipeline if a counter associated with a primary instruction from which the secondary instruction depends indicates that there is a hazard related to that primary instruction.

(9) A second aspect provides a method to avoid data hazards in a parallel processing unit, the method comprising: decoding, by an instruction decoder, an instruction; in response to determining at the instruction decoder that the decoded instruction is a primary instruction from which at least one other instruction is dependent on, causing a value of a counter of a plurality of counters that is associated with the primary instruction to be adjusted to indicate that there is a hazard related to the

primary instruction; forwarding the decoded instruction from the instruction decoder to a queue of a plurality of queues, each queue to receive instructions to be executed by one of a plurality of instruction pipelines; in response to determining, at the queue, that a received instruction is a secondary instruction that is dependent on one or more primary instructions, stalling the secondary instruction from execution by the associated instruction pipeline if a counter associated with a primary instruction from which the secondary instruction depends indicates that there is a hazard related to the primary instruction; and in response to detecting, by monitor hardware logic, that a hazard related to a primary instruction has been resolved by an instruction pipeline of the plurality of instruction pipelines, causing the value of the counter associated with the primary instruction to be adjusted to indicate that the hazard related to the primary instruction has been resolved.

(10) A third aspect provides a computer-implemented method of generating computer executable instructions for a parallel processing unit, the method comprising, by a processor: receiving a plurality of related instructions; identifying data hazards in the plurality of related instructions, each data hazard comprising a primary instruction and one or more secondary instructions; allocating each primary instruction a counter of a plurality of counters for tracking the identified data hazard; generating a computer executable instruction for each primary instruction that comprises information indicating the computer executable instruction is a primary instruction and information identifying the counter allocated to the primary instruction; and generating a computer executable instruction for each secondary instruction that comprises information identifying the computer executable instruction is a secondary instruction and information identifying the counter allocated to the corresponding primary instruction; and loading the computer executable instructions into the parallel processing unit.

(11) The parallel processing units described herein may be embodied in hardware on an integrated circuit. There may be provided a method of manufacturing, at an integrated circuit manufacturing system, the parallel processing units described herein. There may be provided an integrated circuit definition dataset that, when processed in an integrated circuit manufacturing system, configures the system to manufacture the parallel processing units described herein. There may be provided a non-transitory computer readable storage medium having stored thereon a computer readable description of an integrated circuit that, when processed in an integrated circuit manufacturing system, causes the integrated circuit manufacturing system to manufacture the parallel processing units described herein.

(12) There may be provided an integrated circuit manufacturing system comprising: a non-transitory computer readable storage medium having stored thereon a computer readable integrated circuit description that describes the parallel processing units described herein; a layout processing system configured to process the integrated circuit description so as to generate a circuit layout description of an integrated circuit embodying the parallel processing units described herein; and an integrated circuit generation system configured to manufacture the parallel processing units described herein according to the circuit layout description.

(13) There may be provided computer program code for performing a method as described herein. There may be provided non-transitory computer readable storage medium having stored thereon computer readable instructions that, when executed at a computer system, cause the computer system to perform the methods as described herein.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The present invention is described by way of example with reference to the accompanying drawings. In the drawings:

(2) FIG. 1 is a block diagram of a first example parallel processing unit;

- (3) FIG. 2 is a schematic diagram of the format of an instruction of FIG. 1;
- (4) FIG. 3 is a block diagram of an example queue of FIG. 1;
- (5) FIG. 4 is a block diagram of an example counter block of FIG. 1;
- (6) FIG. 5 is a flow diagram of an example method for avoiding inter-pipeline data hazards;
- (7) FIG. 6 is a block diagram of a second example parallel processing unit;
- (8) FIG. 7 is a block diagram of an example counter block of FIG. 6;
- (9) FIG. 8 is a flow diagram of a second example method for avoiding inter-pipeline data hazards;
- (10) FIG. 9 is a block diagram of a third example parallel processing unit;
- (11) FIG. 10 is a schematic diagram of an example counter block of FIG. 9;
- (12) FIG. 11 is a flow diagram of a third example method for avoiding inter-pipeline data hazards;
- (13) FIG. 12 is a flow diagram of an example method of generating instructions for execution on a parallel processing unit;
- (14) FIG. 13 is a flow diagram of an example method of allocating counters to the primary instructions;
- (15) FIG. 14 is a block diagram of an example computer system in which a parallel processing unit described herein is implemented; and
- (16) FIG. 15 is a block diagram of an example integrated circuit manufacturing system for generating an integrated circuit embodying a system to implement the parallel processing units described herein.

DETAILED DESCRIPTION OF THE INVENTION

- (17) The following description is presented by way of example to enable a person skilled in the art to make and use the invention. The present invention is not limited to the embodiments described herein and various modifications to the disclosed embodiments will be apparent to those skilled in the art. Embodiments are described by way of example only.
- (18) As described above, many processing units, such as GPUs, comprise a plurality of parallel instruction pipelines which are designed to efficiently process large amounts of data in parallel. For example, some processing units, may comprise a set of parallel instruction pipelines which include at least two instruction pipelines that are each optimized for a particular type (or types) of computation. Having multiple instruction pipelines that are configured to execute different types of computations allows slow or rarely used instructions to be executed in parallel with high-throughput common arithmetic operations so that the slow, or rarely used, instructions do not become a bottleneck. This also allows the Arithmetic Logic Units (ALUs) of each pipeline to be separately optimised for their particular use.
- (19) While a plurality of instruction pipelines allows for more efficient use of processing resources (e.g. Arithmetic Logic Unit (ALU) resources) and allows stalls caused by resource contention to be hidden, reordering the instructions over multiple instruction pipelines complicates tracking data hazards to ensure that instructions are performed in the correct order.
- (20) In particular, in such parallel processing units not only do intra-pipeline hazards (i.e. hazards related to instructions executed by the same pipeline) need to be tracked and eliminated, but inter-pipeline hazards (i.e. hazards related to instructions executed by different pipelines) also need to be tracked and eliminated. Specifically, since there are multiple instruction pipelines running in parallel related instructions may be in different pipelines (with different processing rates) at the same time. Accordingly, what is needed is a mechanism that ensures that if a data hazard exists between instructions executed in different pipelines that the dependent instruction will not be executed until the data hazard has cleared.
- (21) Detecting inter-pipeline data hazards solely in hardware is very costly in terms of area due to the significant number of pipeline stages that would need to be tracked and the significant number of comparisons that would be required.
- (22) Accordingly, described herein are software-controlled methods and systems for avoiding inter-pipeline data hazards in a GPU or other parallel processing units (such as for high performance

computing applications) with a plurality of parallel instruction pipelines. In particular, in the methods and systems described herein inter-pipeline data hazards are identified at build time (e.g. by a compiler) and information is inserted in the instructions that identifies primary instructions (i.e. instructions from which one or more instructions in another pipeline depends) and secondary instructions (i.e. instructions that depend on one or more primary instructions in another pipeline) and links the primary and secondary instructions via a counter which is used to track the inter-pipeline data hazard and enforce the appropriate ordering of instructions.

(23) When the instruction decoder of the parallel processing unit outputs a primary instruction for execution the associated counter is modified (e.g. incremented) to indicate that there is a hazard related to that primary instruction (i.e. that it is not safe to execute secondary instructions that are dependent on that primary instruction). When it is subsequently detected that the hazard related to that primary instruction has been resolved (e.g. the primary instruction has written data to memory) the value of the associated counter is adjusted (e.g. decremented) to indicate that the hazard related to that primary instruction has been resolved (i.e. that it is safe to execute secondary instructions that are dependent on that primary instruction). Instructions output by the instruction decoder for execution are sent to a queue associated with the appropriate instruction pipeline. Prior to sending an instruction from the queue to the instruction pipeline for execution the queue checks, for each secondary instruction, the counter(s) associated with the primary instruction(s) from which the secondary instruction depends. So long as at least one of the counter(s) associated with a primary instruction(s) from which the secondary instruction depends indicates there is a hazard the secondary instruction is stalled in the queue.

(24) Stalling secondary instructions right before they are to be executed by an instruction pipeline has shown to improve performance in cases where the primary instruction(s) on which the secondary instruction depends will be completed quickly (e.g. when a primary instruction is executed by an instruction pipeline with high throughput). Such inter-pipeline data hazards may be referred to herein as low latency inter-pipeline data hazards. However, stalling secondary instructions right before they are to be executed by an instruction pipeline has shown to reduce performance where the primary instruction(s) on which the secondary instruction depends will be completed slowly (e.g. when a primary instruction is executed by an instruction pipeline with low throughput). Such inter-pipeline data hazards may be referred to herein as high latency inter-pipeline data hazards.

(25) Accordingly, in some embodiments described herein the compiler may be configured to separately identify and mark low latency inter-pipeline data hazards and high latency inter-pipeline data hazards. In these embodiments, the low latency inter-pipeline data hazards may be processed as described above (e.g. when an instruction decoder outputs a primary instruction of a low latency inter-pipeline data hazard the value of a counter associated with the primary instruction is adjusted to indicate there is a hazard related to that primary instruction and when it is subsequently detected that the hazard related to that primary instruction has been resolved (e.g. the primary instruction has written data to memory) the value of the counter associated with the primary instruction is adjusted to indicate that the hazard related to that primary instruction has been resolved; and secondary instructions related to a low latency data hazard that have been output by the instruction decoder for execution are stalled in a queue preceding the appropriate instruction pipeline so long as the value of at least one of the counters associated with the primary instructions from which it depends indicate that there is a hazard).

(26) The high latency inter-pipeline data hazards, however, are processed in a different manner. Specifically, the primary instructions of high latency inter-pipeline data hazards are processed in the same manner as the primary instructions of low latency inter-pipeline data hazards (e.g. when a primary instruction of a high latency inter-pipeline data hazard is output by a decoder for execution by an instruction pipeline the value of a counter associated therewith is adjusted to indicate there is a data hazard related to the primary instruction and when it is subsequently detected that the hazard

has been resolved (e.g. the primary instruction has written to memory) the value of the counter associated therewith is adjusted to indicate that the data hazard related to the primary instruction has been resolved). However, when secondary instructions of at least one high latency data hazard are decoded by the instruction decoder a determination is made then as to whether the relevant high latency data hazard(s) have been resolved (i.e. whether the values of the counters associated with the primary instruction(s) from which it depends indicate that high latency hazard has been resolved). If the relevant high latency hazards have been resolved, the secondary instruction is output by the decoder for execution by the appropriate instruction pipeline. If, however, at least one relevant high latency hazard has not been resolved, then the instruction decoder de-schedules the secondary instruction (e.g. sends the secondary instruction back to a scheduler) until the relevant high latency hazards have been resolved (i.e. until the counters associated with the primary instructions from which it depends indicate that the hazard has been resolved). Once the relevant high latency hazards for the secondary instruction have been resolved the secondary instruction is rescheduled and sent back to the instruction decoder for processing.

(27) In some cases, a secondary instruction may be dependent on both a high latency primary instruction and a low latency primary instruction. In these cases, the secondary instruction would be subject to both inter-pipeline hazard avoidance mechanisms described above. Specifically, the instruction decoder would check the counters associated with the high latency primary instructions and the queue would be configured to check the counters associated with the low latency primary instructions.

(28) While the methods, systems and techniques described herein are described as being used for inter-pipeline data hazard avoidance, the methods, systems and techniques described herein may also be used for intra-pipeline data hazards. For example, the methods, systems and techniques described herein may be also be used for intra-pipeline data hazard avoidance in cases where the area versus performance trade-off does not justify the cost of having cycle-accurate hazard detection which may be achieved by other methods. In these cases, the compiler would be configured to also identify intra-pipeline data hazards and update the primary and secondary instructions thereof in the same manner as described herein.

(29) Reference is now made to FIG. 1 which illustrates a first example parallel processing unit **100** which may be a GPU or other parallel processing unit. It will be appreciated that FIG. 1 only shows some elements of the parallel processing unit **100** and there may be many other elements (e.g. caches, interfaces, etc.) within the parallel processing unit that are not shown in FIG. 1. The parallel processing unit **100** of FIG. 1 comprises a counter block **102** comprising a plurality of counters **103**, an instruction decoder **104**, a plurality of instruction pipelines **106**, monitor logic **108** and a queue **110** preceding each instruction pipeline **106**.

(30) The counter block **102** comprises a plurality of counters **103** that are used to track inter-pipeline data hazards and enforce ordering of the instructions in accordance therewith. In particular, the counters **103** are used to indicate (i) when there is a hazard related to a primary instruction and thus it is not safe for a secondary instruction that is dependent thereon to be executed (e.g. the secondary instruction(s) should stall); and (ii) when the hazard related to a primary instruction has been resolved and thus it is safe for the secondary instruction(s) that are dependent thereon to be executed. Specifically, the counters **103** are configured so that when a counter has predetermined value or set of values it indicates that there is a hazard related to the associated primary instruction; and when a counter has a different predetermined value or set of predetermined values it indicates that the hazard related to the associated primary instruction has been resolved. In some examples, the counters **103** are configured so that when a counter has a non-zero value it indicates that there is a hazard related to the associated primary instruction and when a counter has a zero value it indicates that the hazard related to the associated primary instruction has been resolved. It will be evident to a person of skill in the art that this is an example only and that the counters **103** may be configured so that different values indicate that there is a hazard and/or the hazard has been

resolved.

(31) When a counter **103** indicates that there is a hazard related to the associated primary instruction the counter **103** acts as a fence at which a secondary instruction from which it depends has reached. Specifically, the secondary instruction must wait until the fence is removed. Accordingly, the counters **103** of FIG. **1** may be referred to herein as fence counters.

(32) The counter block **102** is configured to adjust the values of the counters **103** in response to receiving adjustment instructions or signals from the instruction decoder **104**, monitor logic **108**, and optionally the queues **110**; and to generate and provide counter status information to the queues **110**. In particular, as described in more detail below, the instruction decoder **104** is configured to, in response to outputting a primary instruction for execution (e.g. in response to forwarding a primary instruction to a queue **110**) send an adjustment instruction or signal to the counter block **102** that causes the counter block **102** to adjust the value of a counter **103** associated with the primary instruction to indicate there is a hazard with the primary instruction. The monitor logic **108** is configured to, in response to detecting that a hazard related to a primary instruction has been (partially or fully) resolved by an instruction pipeline **106**, send an adjustment instruction or signal to the counter block **102** that causes the counter block **102** to adjust the value of the counter **103** associated with the primary instruction to indicate that the hazard related to the primary instruction has been (partially or fully) resolved. The queue **110** may also be configured to, in response to detecting that a primary instruction is (partially or fully) no longer active and thus the primary instruction is (partially or fully) discarded, send an adjustment instruction or signal to the counter block **102** that causes the counter block **102** to adjust the value of the counter **103** associated with the primary instruction to indicate that the hazard related to the primary instruction has been (partially or fully) resolved.

(33) The counter status information comprises information that indicates whether there is a hazard related to the primary instruction associated with each counter or whether the hazard has cleared, or has been resolved. The counter block **102** is configured to generate the counter status information based on the value of the counters **103**. In some cases, the counter status information may comprise a flag or bit for each counter **103** indicating whether there is a hazard related to the corresponding primary instruction or whether the hazard related to the corresponding primary instruction has been resolved. For example, the counter status information may comprise a single-bit flag for each counter **103** where a flag is set to “1” to indicate that there is a hazard related to the primary instruction and a flag is set to “0” to indicate that the hazard related to the primary instruction has been resolved. In other cases, the counter status information may comprise the value of each of the counters **103** and the recipient of the counter status information is configured to determine from the values whether the hazards related to the associated primary instructions have been resolved.

(34) The instruction decoder **104** receives instructions which include information (inserted at build time—e.g. by a compiler) that identify primary instructions (i.e. instructions from which at least one other instruction in another instruction pipeline is dependent on), secondary instructions (i.e. instructions that are dependent on at least one primary instruction in another pipeline) and the counter(s) they are associated with. Specifically, each primary instruction will be allocated a counter and the secondary instruction(s) will be linked to the primary instruction via that counter. Since a secondary instruction may be dependent on more than one primary instruction, secondary instructions may be linked to multiple primary instructions via multiple counters. An example of the information and format of the information that identifies primary and secondary instructions and the counters they are associated with is described below with reference to FIG. **2**. There are typically fewer counters than there are inter-pipeline data hazards so the counters are generally re-used for multiple inter-pipeline data hazards.

(35) The instruction decoder **104** decodes the received instructions, selects the appropriate instruction pipeline for executing each instruction, and outputs the instructions for execution by the selected instruction pipelines. If the instruction decoder **104** determines that an instruction output

for execution is a primary instruction the instruction decoder sends an adjustment instruction or signal to the counter block **102** that causes the counter block **102** to adjust the value of the counter **103** associated with that primary instruction to indicate that there is a hazard related to the primary instruction (and thus it is not safe to execute secondary instructions that are dependent on that primary instruction). For example, if primary instruction X is associated with counter 2 then when the instruction decoder **104** outputs primary instruction X for execution the instruction decoder **104** will output an adjustment signal or instruction to the counter block **102** that causes the counter block **102** to adjust the value of the counter 2 to indicate that there is a hazard related to primary instruction X.

(36) In some examples, the instruction decoder **104** may be configured to, in response to outputting a primary instruction for execution, output an adjustment instruction or signal that causes the counter block **102** to increment the counter **103** associated with the primary instruction by a predetermined amount (e.g. **8**). In some cases, as described below, each instruction may be part of, or related to, a task that causes multiple instances of the instruction to be executed. In these cases, it may only be safe for a secondary instruction to be executed if the hazard has been resolved for all instances of the primary instruction. In such cases the predetermined amount by which the counter **103** is incremented may reflect the number of instances, or groups of instances, for which the hazard can be separately tracked. For example, if 32 instances of the instruction may be executed in groups of 4 (i.e. 8 groups) the instruction decoder **104** may be configured to increment the counter **103** by 8.

(37) Each instruction pipeline **106** comprises hardware logic (e.g. one or more ALUs) for executing instructions. In some examples, the plurality of instruction pipelines **106** includes at least two different instruction pipelines that are configured to execute decoded instructions of different types. For example, the instruction pipelines **106** may comprise one or more instruction pipelines that are configured to: (i) perform bit integer operations, floating point operations and logical (bitwise) operations; (ii) calculate per-instance texture coordinate or other varyings; (iii) perform 32-bit float non-rational/transcendental operations; (iv) execute 64-bit float operations; (v) perform data copying and format conversion; (vi) execute texture address calculation; and (vii) execute atomic operations on local memory registers. Having multiple instruction pipelines that are configured to execute different types of instructions allows slow or rarely used instructions to be executed in parallel with high-throughput common arithmetic operations so that the slow or rarely used instructions do not become a bottleneck. This also allows ALUs to be separately optimised for their particular use.

(38) In some cases, the instruction pipelines **106** may each be single-instruction multiple-data (SIMD) pipelines. As is known to those of skill in the art, a SIMD instruction is an instruction that, when executed, causes the same operation(s) to be performed on multiple data items that are associated with the instruction. SIMD instructions allow fewer instructions to specify the same amount of work reducing the pressure on the instruction fetch module and the instruction decoder. A SIMD pipeline is thus a pipeline that is able to process SIMD instructions—i.e. it is a pipeline that is able to execute the same instruction on multiple data items. This means that where the instructions are part of tasks, as described in more detail below, the instruction pipelines **106** can execute an entire task's worth of instances or data-items using one issued instruction. The instruction pipeline may take more than one clock cycle to process the issued SIMD instruction.

(39) The monitor logic **108** monitors the instruction pipelines **106** to detect when a hazard related to a primary instruction has been resolved (partially or fully) by an instruction pipeline and in response to detecting that a hazard related to a primary instruction has been (partially or fully) resolved by an instruction pipeline **106** sends an adjustment instruction or signal to the counter block **102** to cause the counter associated with the primary instruction to indicate that the hazard related to the primary instruction has been (partially or fully) resolved. For example, if a primary instruction is associated with counter 2, when the monitor logic **108** detects that an instruction

pipeline **106** has resolved the hazard related to that primary instruction the monitor logic **108** will send an adjustment signal or instruction to the counter block **102** to cause the counter block **102** to adjust the value of the counter 2 to indicate that the hazard associated with that primary instruction has been resolved.

(40) In some examples, the monitor logic **108** may be configured to, in response to detecting that a hazard associated with a primary instruction has been (partially or fully) resolved by an instruction pipeline **106**, send an adjustment signal or instruction to the counter block **102** that causes the counter block **102** to decrement the value of the counter **103** by a predetermined amount (e.g. 1 or 8) to indicate that the hazard has been (partially or fully) resolved.

(41) As described in more detail below, in some cases, each instruction may be part of, or associated with, a task which causes multiple instances (e.g. up to 32 instances) of the instruction to be executed. In these cases, the hazard is said to be fully resolved when the hazard has been resolved by all instances, and the hazard is said to be partially resolved when the hazard has been resolved by some (but not all) of the instances. The instances may be divided into a number of groups (e.g. 8) and each group is executed as a block such that the execution of each block can be tracked separately. In these cases, the monitor logic **108** may be configured to send a separate instruction or signal each time it detects that a hazard related to a primary instruction has been resolved by a group of instances to cause the value of the counter to be adjusted to indicate that the hazard has been partially resolved (e.g. an instruction to decrement the value of the counter by 1). Once the hazard is resolved by each group the counter will indicate that the hazard has been fully resolved. It will be evident to a person of skill in the art that this is an example only and that the monitor logic **108** may be configured to cause the counter block **102** to adjust the value of the counter associated with a primary instruction in any suitable manner so that the counter will have a value indicating that the hazard related thereto has been (fully or partially) resolved.

(42) The monitor logic **108** may be configured to use different criteria to determine when a hazard has been resolved by an instruction pipeline **106** based on the type of hazard. For example, a WAW or a RAW hazard may be resolved when the primary instruction has written the result of the instruction to storage **112** such as memory **114** or a register (not shown). Accordingly, the monitor logic **108** may be configured to detect that a WAW or RAW hazard has been resolved by an instruction pipeline **106** when the monitor logic **108** detects that an instruction pipeline **106** has written the result of a primary instruction to storage **112**. In these cases, where each instruction pipeline **106** has an interface to the storage units, the monitor logic **106** may be configured to monitor these instruction pipeline **106** to storage interfaces to detect writes to the storage. In contrast, a WAR hazard may be resolved when the sources for the primary instruction have been read by the instruction pipeline **106**. Accordingly, the monitor logic **108** may be configured to detect that a WAR hazard has been resolved by an instruction pipeline **106** when the monitor logic **108** detects that the sources for a primary instruction have been read by an instruction pipeline **106**.

(43) Although the monitor logic **108** is shown in FIG. 1 as being a single logic block that is separate from the instruction pipelines, in other examples the monitor logic **108** may be distributed amongst, and part of, the instruction pipelines **106**. For example, each instruction pipeline **106** may comprise its own monitor logic.

(44) Each instruction pipeline **106** is preceded by a queue **110** that receives instructions from the instruction decoder **104** that are to be executed by the corresponding instruction pipeline **106** and forwards the received instructions to the corresponding instruction pipeline **106** for execution in order. Each queue **110** is configured to, prior to forwarding an instruction to the corresponding instruction pipeline **106** for execution, determine whether the instruction is a secondary instruction. If the instruction is not a secondary instruction then the instruction is forwarded to the corresponding instruction pipeline **106** for execution. If, however, the instruction is a secondary instruction then a determination is made (from the counters **103** and/or the counter status information) whether the hazards related to the primary instructions from which the secondary

instruction depends have been resolved. If the hazards related to the primary instruction from which the secondary instruction depends have been resolved then the secondary instruction is forwarded to the corresponding instruction pipeline **106** for execution. If, however, at least one of the hazards related to a primary instruction from which the secondary instruction depends have not been resolved then the instruction is stalled.

(45) Accordingly, only if all the counters associated with the primary instruction(s) indicate that the related hazard has been resolved can an instruction be forwarded to the instruction pipeline **106** for execution. For example, if a queue **110** receives a secondary instruction that is dependent on the primary instructions associated with counters 2 and 3 then the queue **110** cannot forward the secondary instruction to the instruction pipeline until counters 2 and 3 both have a value (e.g. zero) indicating that the related hazards have been resolved.

(46) In some examples, stalling a secondary instruction may stall all subsequent instructions from being executed by the associated instruction pipeline. However, as described in more detail below with reference to FIG. 3, in other examples, where the instructions are part of, or associated with tasks, while a queue **110** stalls a secondary instruction related to a first task it may be able to forward other later instructions related to a different task to the associated instruction pipeline.

(47) In some cases, the queue **110** may be configured to determine the value of the appropriate counter(s) by polling or requesting counter status information for the appropriate counters from the counter block **102**. In other cases, the counter block **102** may be configured to periodically push the counter status information to the queues **110**.

(48) In some cases, the queue **110** may also be configured send an adjustment instruction or signal to the counter block **102** that causes the counter block **102** to adjust the value of the counter **103** associated with the primary instruction to indicate that the hazard has been (partially or fully) resolved if the queue **110** detects, prior to forwarding, that the instruction is to be (partially or fully) discarded. Specifically, for a variety of reasons it may be possible for an instruction to be sent to the queue **110** for execution, but when it is time for that instruction to be issued to the instruction pipeline it may no longer be desirable for that instruction to be executed.

(49) For example, this may occur when the parallel processing unit implements predication. As is known to those of skill in the art, predication is a process implemented in parallel processing units that is an alternative to branch prediction. In branch prediction the parallel processing unit predicts the path of a branch that will be executed and predictively executes the instructions related to that branch. A mis-prediction (i.e. an incorrect guess of which path of the branch will be taken) can result in a stall of a pipeline and cause instructions to be fetched from the actual branch target address. In contrast, in predication instructions related to all possible paths of a branch are executed in parallel and only those instructions associated with the taken path (as determined from the branch condition) are permitted to modify the architecture state. Each instruction from a particular path will be associated with a predicate (e.g. Boolean value) which indicates whether the instruction is allowed to modify the architecture state or not. The predicate value will be set based on the evaluation of the branch condition. An instruction whose predicate indicates that the instruction is not allowed to modify the architecture state is said to have been predicated out. If an instruction has been predicated out before it is forwarded to an instruction pipeline then there is no need to forward it to the instruction pipeline for execution.

(50) Accordingly, before an instruction is forwarded to the instruction pipeline for execution the queue **110** may be configured to determine based on active information (e.g. predicate information) whether it is desirable to forward the instruction to the pipeline for execution. If the active information indicates that the instruction is not to be executed (e.g. the predicate indicates the instruction has been predicated out) then the instruction is discarded, and, if the instruction is a secondary instruction, the queue **110** sends an adjustment instruction or signal to the counter block **102** that causes the counter block **102** to adjust the value of the counter **103** associated with that primary instruction to indicate that the hazard has been resolved (e.g. an instruction that causes the

counter block to decrement the value of the counter by a predetermined amount).

(51) Where the instructions are associated with tasks then it is possible for the instruction to be active for some instances and not others. This may occur for example, where the instructions are predicated on a per instance basis. In these cases, the queue **110** may be configured to detect if the instruction is partially active (some but not all instances are active), fully inactive (all instance are inactive), or fully active (all instances are active). If the queue **110** detects that the instruction is fully inactive the queue **110** may send an adjustment instruction to the counter block **102** that causes the counter block **102** to adjust the value of the counter **103** associated with that primary instruction to indicate that the hazard has been fully resolved (e.g. an instruction that causes the counter block to decrement the value of the counter by 8), and if the queue **110** detects that the instruction is partially inactive the queue **110** may send an adjustment instruction to the counter block **102** that causes the counter block **102** to adjust the value of the counter **103** associated with the primary instruction to indicate that the hazard has been partially resolved (e.g. an instruction that causes the counter block to decrement the value of the counter by less than 8 to reflect what portion of the instruction the hazard has been resolved).

(52) An example implementation of the queue **110** of FIG. 1 is described below with reference to FIG. 3.

(53) In some examples, the instructions decoded by the instruction decoder **104** and executed by the instruction pipelines **106** each relate to or correspond to a task. The term ‘task’ is used herein to refer to a group or body of work (i.e. a plurality of work-items or program instances) in which the same program (e.g. the same sequence of ALU instructions) is applied to a set of data, where this data may comprise one or more elements (or data-items, e.g. a plurality of pixels or vertices). A task therefore refers to one or more data-items (one or more references/points to data-items) and a program (or reference (e.g. pointer) to a program (i.e. a sequence of instructions)) which are to be executed on the data-items. As described above, when an instruction related to a task is executed an instance of that instruction is independently executed against each of the data items.

(54) In these examples, tasks may be received at a scheduler **116** which schedules the received tasks for execution (i.e. determines the order in which the tasks are to be executed). Once the scheduler **116** determines which task is to be executed next, the scheduler **116** sends information identifying that task (e.g. a task ID) and information identifying the next instruction to be fetched (e.g. program counter value) to an instruction fetch module **120**. The instruction fetch module **120** then fetches the next instruction for that task from memory (e.g. via a cache structure) based on the received information (e.g. the program counter value). The fetched instruction is then sent to the instruction decoder **104** to be decoded. A task that has been sent to the instruction fetch module **120** for execution is said to be an active task or a scheduled task. The instruction fetch module **120** and the instruction decoder **104** may form part of an instruction controller **118**.

(55) Since tasks contain related instructions (e.g. instructions of the same program) there are typically dependencies (e.g. hazards) between instructions within a task, but there are not typically dependencies (e.g. hazards) between instructions in different tasks. As a result, there may be a plurality of counters **103** associated with each task ID to track the inter-pipelines hazards within a task via that task ID. An example of a counter block **102** wherein the counters **103** are associated with one of the task IDs is described below in relation to FIG. 4. Where the counters associated with each task ID are assigned the same counter IDs (e.g. counter numbers) the specific counter associated with an instruction may be identified or indexed using both the task ID and the counter ID.

(56) Reference is now made to FIG. 2 which illustrates an example format of an instruction **200** in the system of FIG. 1 that comprises information identifying primary instructions and secondary instructions and the counters they are associated with. In the example of FIG. 2 the instruction comprises a primary instruction field **202** and a secondary instruction field **204**. The primary instruction field **202** indicates whether or not the instruction is a primary instruction and if so what

counter the primary instruction is associated with. In particular, in the example of FIG. 2 the primary instruction field **202** is configured to hold an N-bit number. When the N-bit number has a particular value (e.g. zero) it indicates that the instruction is not a primary instruction, and when the N-bit number has a value other than the particular value (e.g. a non-zero value) it indicates that the instruction is a primary instruction and it identifies the counter that the primary instruction is associated with (e.g. by number).

(57) For example, a three-bit primary instruction field **202** with a binary value of “000” (i.e. a decimal value of zero) may indicate that the instruction is not a primary instruction (i.e. no other instructions in another instruction pipeline are dependent on this instruction) and a three-bit primary instruction field **202** with a binary value of “001” (i.e. a decimal value of one) may indicate that the instruction is a primary instruction (i.e. there is at least one other instruction in another instruction pipeline that is dependent on this instruction) and that the primary instruction is associated with counter number **1**.

(58) The number of bits (i.e. N) allocated to the primary instruction field **202** will be based on the number of counters **103**. Specifically, the primary instruction field **202** generally comprises enough bits to uniquely identify each of the counters. For example, if there are six counters then the primary instruction field may comprise 3 bits (which can represent 8 unique numbers) to be able to uniquely identify the six counters by a unique number.

(59) The secondary instruction field **204** indicates whether the instruction is a secondary instruction and if so, which counters are associated with the primary instructions it is dependent on. In the example of FIG. 2 the secondary instruction field **204** is configured to hold an M-bit mask. There is one bit per counter and the value of that bit indicates whether the instruction is dependent on the primary instruction associated with that counter. For example, if there are six counters there will be six bits in the mask, the first bit may correspond to the first counter, the second bit may correspond to the second counter, the third bit may correspond to the third counter etc. In some cases, a bit that is set (e.g. is “1”) may indicate that the instruction is dependent on the primary instruction associated with the corresponding counter, and a bit that is cleared, or is not set, (e.g. is “0”) may indicate that the instruction is not dependent on the primary instruction associated with the corresponding counter. In these cases, if at least one of the bits in the mask of the secondary instruction field **204** is set then the instruction is dependent on at least one instruction in another pipeline and thus the instruction is a secondary instruction; and if none of the bits in the mask of the secondary instruction field **204** are set then the instruction is not dependent on any instructions in another pipeline and thus the instruction is not a secondary instruction.

(60) For example, a 6-bit secondary instruction field **204** with a binary value of “000000” may indicate that the instruction is not a secondary instruction (i.e. it is not dependent on any other instructions) and a 6-bit secondary instruction field **204** with a binary value of “010100” may indicate that the instruction is a secondary instruction that is dependent on the instructions associated with the second and fourth counters.

(61) Since an instruction can be both dependent on one or more other instructions and have one or more other instructions depend on it, it is possible for an instruction to be both a primary instruction and a secondary instruction.

(62) As described above, the instructions, such as instructions **200** of FIG. 2, may be generated at build time by a compiler that is configured to identify inter-pipeline data hazards and allocate a counter to each identified inter-pipeline data hazard to be used to track the inter-pipeline data hazard and enforce the appropriate order. The primary instruction of that data hazard is then modified to include information (e.g. in the primary instruction field **202**) that indicates that it is a primary instruction and the counter that has been associated with that primary instruction; and the secondary instruction(s) of the data hazard are modified to include information (e.g. in the secondary instruction field **204**) that indicates the counters associated with the primary instructions that it is dependent on.

(63) In some cases, the compiler may be configured to allocate one counter to each identified inter-pipeline data hazard. In other cases, the compiler may be configured to use a single counter for multiple identified inter-pipeline data hazards. For example, if instructions C and D are to be executed in the same instruction pipeline and C is dependent on instruction A and instruction D is dependent on instruction B then primary instructions A and B may both be allocated the same counter. C will then have to wait for the hazard associated with instructions A and B to be resolved before it can be executed and since D follows C will also have to wait for the hazard associated with instruction A and B to be resolved. If instructions C and D are to be executed in different instruction pipelines then they will both wait for the hazards associated with instruction A and B to be resolved, but they may have waited too long. Using a single counter for multiple inter-pipeline data hazards are preferred where a secondary instruction is dependent on two different primary instructions in different pipelines. In such cases the primary instructions could both be allocated the same counter. Accordingly, in some cases there may be a performance cost in using a single counter for multiple inter-pipeline data hazards, however this may allow a smaller number of counters to be used thereby decreasing the size of the hardware implementation.

(64) Reference is now made to FIG. 3 which illustrates an example implementation of a queue **110** of FIG. 1 for use in a parallel processing unit wherein each instruction corresponds to a task that has been allocated a task ID and each instruction pipeline **106** receives instructions relating to multiple tasks each with a corresponding different task ID. As described above, the queue **110** receives instructions from the instruction decoder **104** to be executed by the corresponding instruction pipeline **106** and forwards the instructions to the corresponding instruction pipeline **106** in the correct order. However, any secondary instruction cannot be forwarded to the instruction pipeline **106** until the counters associated with the primary instructions from which the secondary instruction depends indicate the hazards related to the primary instructions have been resolved. In this example, the queue **110** preferably issues all instructions to the corresponding instruction pipeline in the order in which they were received, but the queue **110** must maintain the order of the instructions within a task. This means that if a secondary instruction is stalled because the counters associated with the primary instruction(s) from which the secondary instruction depends do not indicate that the hazard has been resolved then only the subsequent instructions related to the same task (e.g. identified by the same task ID) also need to be stalled and subsequent instructions related to a different task (e.g. identified by a different task ID) may be forwarded to the corresponding instruction pipeline **106** before the stalled secondary instruction.

(65) In the example of FIG. 3 the queue **110** comprises a FIFO (first in first out) queue **302**, hazard detection logic **304**, a hazard instruction queue **306**, an information queue **308** and a MUX (multiplexer) **310**.

(66) The FIFO queue **302** receives instructions output by the instruction decoder **104**, stores the instructions in the order they were received, and then outputs them to the hazard detection logic **304** in the same order (i.e. the order in which they were received). In some cases (e.g. when there is one instruction decoder), the FIFO queue **302** may be capable of receiving a maximum of one instruction per clock cycle. In other cases (e.g. when there are multiple instruction decoders), the FIFO queue **302** may be capable of receiving more than one instruction per clock cycle. The FIFO queue **302** can typically only output a maximum of one instruction per clock cycle.

(67) The hazard detection logic **304** is configured to determine whether the instruction output from the FIFO queue **302** is hazard free or not (i.e. whether the instruction can be safely executed). In this example, there are two potential hazards: (i) the instruction may be a secondary instruction that is waiting for the hazard(s) related to one or more primary instructions to be resolved before it can be issued to the instruction pipeline for execution; and/or (ii) the instruction may relate to the same task (e.g. has the same task-ID) as an instruction in the hazard instruction queue **306** and thus has to wait for that instruction to be issued to the corresponding instruction pipeline before it can be issued to the instruction pipeline for execution. Since a younger instruction related to a task cannot

be issued to the instruction pipeline before an older instruction related to the same task, if there is already an instruction in the hazard instruction queue **306** related to the same task that has not been issued to the instruction pipeline then the new instruction must wait for the older instruction to be issued to the instruction pipeline before it can be issued to the instruction pipeline.

(68) In some cases, the hazard detection logic **304** is configured to first determine whether there are any instructions in the hazard instruction queue **306** that are related to the same task (e.g. based on the task ID associated with the instruction). If it is determined that there is at least one stalled instruction in the hazard instruction queue **306** with the same task ID then the instruction cannot be safely executed and the instruction is sent to the hazard instruction queue **306** where it waits for the earlier instruction with the same task ID to be sent to the corresponding instruction pipeline **106**.

(69) If, however, there are no instructions in the hazard instruction queue **306** that have the same task ID then it is determined whether the instruction received from the FIFO queue **302** is a secondary instruction (indicating a potential hazard). If the hazard detection logic **304** determines that the instruction received from the FIFO queue **302** is not a secondary instruction then the instruction is a hazard free instruction (and thus is ready to be executed) and the instruction is forwarded directly to the MUX **310**. If, however, the hazard detection logic **304** determines that the instruction received from the FIFO queue **302** is a secondary instruction (e.g. the secondary instruction field of the instruction is non-zero) then the hazard detection logic **304** determines whether the counters associated with the primary instructions from which it depends indicate that the hazard related to the associated primary instruction has been resolved (e.g. via the counter status information).

(70) If all relevant counters indicate that the hazard related to the associated primary instructions have been resolved (e.g. all the relevant counters are zero) then the instruction is a hazard free instruction (and thus is ready for execution) and the instruction is forwarded directly to the MUX **310**. If, however, at least one of the relevant counters indicates that the hazard related to the associated primary instruction has not been resolved (e.g. at least one of the relevant counters has a non-zero value) then the instruction is a hazard instruction and is forwarded to the hazard instruction queue **306**.

(71) The hazard instruction queue **306** is configured to store hazard instructions in order until the corresponding hazard(s) have cleared. The information queue **308** stores information for each hazard instruction in the hazard instruction queue **306** indicating the hazard(s) associated with that instruction. For example, the information queue **308** may store, for each hazard instruction, information indicating whether the hazards related to each primary instruction from which it depends have been resolved and/or information indicating whether the instruction is waiting for another instruction in the hazard queue and if so, which one.

(72) As described above, in this example there are two potential hazards: (i) the instruction may be a secondary instruction that is waiting on the hazard related to one or more primary instructions to be resolved; and/or (ii) the instruction may relate to the same task-ID as an instruction in the hazard instruction queue **306** and is waiting for that instruction to be issued to the corresponding instruction pipeline. Accordingly, the information queue **308** may store a single bit or flag indicating whether the corresponding hazard instruction is waiting for a hazard related to a primary instruction to be resolved (which may be set if at least one of the counters associated with a primary instruction from which the instruction depends indicates that the hazard has not been resolved and may be cleared once all of the counters associated with the primary instruction from which the instruction depends indicates that the hazards have been resolved); and/or the information queue **308** may store a flag or bit indicating whether or not the instruction is waiting for another instruction in the hazard instruction queue **306** to be issued to the corresponding instruction pipeline.

(73) When the information in the information queue **308** indicates that all hazards have cleared (e.g. the counter values for the relevant counters indicate that the hazard related to the primary

instructions from which it depends have been resolved and/or the hazard instruction that it is waiting for has been forwarded to the corresponding instruction pipeline) for a particular instruction then the hazard instruction queue **306** forwards the cleared hazard instruction to the MUX **310**.

(74) When the corresponding instruction pipeline **106** is ready to receive the next instruction the MUX **310** is configured to select one of: the instruction output by the hazard detection logic **304**; and the instruction output by the hazard instruction queue **306**, to forward to the corresponding instruction pipeline **106** for execution. If only one of the hazard detection logic **304** and the hazard instruction queue **306** has output an instruction in the same clock cycle then that MUX **310** selects and forwards that instruction to the corresponding instruction pipeline. If, however, both the hazard detection logic **304** and the hazard instruction queue **306** output an instruction in the same clock cycle then the MUX **310** selects the instruction output from the hazard instruction queue **306** as typically this instruction is older than the instruction output by the hazard detection logic **304**.

(75) In some cases, the queue **110** may also comprise active logic **312** configured to receive the instruction output by the MUX **310** and receive active information that indicates whether all or part of the instruction is active. As described above, it may be possible for an instruction to be sent to the queue for execution, but when it is time for that instruction to be issued to the instruction pipeline it may no longer be desirable for that instruction to be executed (e.g. it may have been predicated out). Accordingly, before the instruction output by the MUX **310** is forwarded to the instruction pipeline for execution the active logic **312** is configured to determine whether it is desirable to forward the instruction to the pipeline for execution based on active information. If the active information indicates that the instruction is not to be executed then the instruction is discarded, and if the instruction is a primary instruction the active logic **312** sends an instruction to the counter block **102** that causes the counter block **102** to adjust the value of the counter **103** associated with that primary instruction to indicate that the hazard has cleared (e.g. an instruction that causes the counter block to decrement the value of the counter by a predetermined amount).

(76) Where the instructions are associated with tasks then it is possible for the instruction to be active for some instances and not others. Where the instances are divided into groups which can be individually tracked the active logic **312** may be configured to determine which groups of instances are active (i.e. have at least one active instance). If there are no active groups for an instruction then the instruction is fully inactive and the instruction is discarded. If the fully inactive instruction is a primary instruction the active logic **312** sends an instruction to the counter block **102** that causes the counter block **102** to adjust the value of the counter **103** associated with that primary instruction to indicate that the hazard has been resolved (e.g. an instruction that causes the counter block to decrement the value of the counter by a predetermined amount) since the hazard will not be resolved or cleared by the instruction pipeline **106**.

(77) If, however, there is at least one active group for the instruction then the instruction is partially inactive (or partially active) and the active logic **312** forwards the instruction to the instruction pipeline **106** for execution. If the partially inactive instruction is a primary instruction the active logic **312** sends an instruction to the counter block **102** that causes the counter block **102** to adjust the value of the counter to indicate that the hazard has been partially resolved (e.g. an instruction that cause the counter block to decrement the value of the counter by the number of inactive groups) since the instances for the inactive group(s) will not be executed and thus the hazard will not be resolved or cleared by the instruction pipeline **106** for these groups.

(78) Reference is now made to FIG. **4** which illustrates an example implementation of the counter block **102** of FIG. **1** wherein each task ID is associated with a plurality of counters **103**. The counter block **102** of FIG. **4** comprises a matrix **402** of counters **103**, counter update logic **404**, and status logic **406**.

(79) As described above, instructions that are associated with or form part of a task are related instructions and thus it is possible that inter-pipeline data hazards may occur between instructions

within a task. However, since tasks themselves are generally independent of other tasks inter-pipeline data hazards will likely not occur between instructions of different tasks. Accordingly, inter-pipeline data hazards can typically be identified and tracked on a per task basis.

(80) The matrix **402** of counters **103** comprises a row for each possible task ID. For example, where there are 3 possible task IDs the matrix **402** will have three rows as shown in FIG. 4. Each row comprises K counters associated with the corresponding task ID where K is an integer greater than or equal to 2. Accordingly, the counters **103** in a particular row are used to track inter-pipeline data hazards for tasks with the corresponding task ID. For example, in FIG. 4 the counters in the first row of the matrix **402** are used to track inter-pipeline data hazards for tasks with a task ID of “01” (binary) and the second row of the matrix **402** are used to track inter-pipeline data hazards for tasks with a task ID of “10” (binary).

(81) As described above, at build time the inter-pipeline data hazards in a task are identified and the primary instructions are amended to include information identifying the instruction as a primary instruction and information identifying the counter to be used to track the hazard (e.g. counter number); and the secondary instructions are amended to include information identifying the instruction as a secondary instruction and information identifying the counter(s) that is/are used to track the primary instruction(s) from which it depends. Accordingly, in addition to identifying inter-pipeline data hazards in a task and identifying the primary and secondary instructions thereof, the compiler is configured to allocate a counter to each identified inter-pipeline data hazard which is used to track the corresponding data hazard. The compiler may be configured to allocate the counters (e.g. counter numbers) to identified inter-pipeline data hazards in any suitable manner. Where there are a small number of counters available per task so that a counter is likely be used more than once for the same task the compiler may be configured to allocate the counters to the identified inter-pipeline data hazards in a manner that maximizes the time between uses of that counter (e.g. the compiler may be configured to allocate the counters to identified inter-pipeline data hazards on a round-robin basis).

(82) The counter update logic **404** receives adjustment instructions or signals from the instruction decoder **104**, the monitor logic **108** and, optionally the queues **110**, which cause the value of a counter to be adjusted to reflect the current status of the corresponding inter-pipeline data hazard. In particular, as described above, when the instruction decoder outputs a primary instruction for execution the instruction decoder **104** sends an adjustment instruction to the counter block **102** to cause the counter **103** associated with the primary instruction to have a value that indicates that there is a hazard related to that primary instruction. In some examples, the instruction decoder **104** is configured to transmit an adjustment instruction to the counter block **102** that causes the value of the counter **103** associated with the primary instruction to be incremented by a predetermined amount (e.g. 8).

(83) When the monitor logic **108** subsequently detects that a hazard related to a primary instruction has been (fully or partially) resolved by one of the instruction pipelines (e.g. by monitoring the interfaces between the instruction pipelines and the storage modules) the monitor logic **108** sends an adjustment instruction or signal to the counter block **102** that causes the counter **103** associated with the primary instruction to indicate that the hazard related to the primary instruction has been (partially or fully) resolved. In some examples, the monitor logic **108** is configured to transmit an adjustment instruction to the counter block **102** that causes the value of the counter **103** associated with the primary instruction to be decremented by a predetermined amount (e.g. 8).

(84) If the queue **110** detects that a primary instruction to be forwarded to the corresponding pipeline is (fully or partially) inactive the queue **110** sends an adjustment instruction or signal to the counter block **102** that causes the counter block **102** to adjust the value of the counter **103** associated with the primary instruction to be adjusted to indicate that the hazard related to the primary instruction has been (partially or fully) resolved.

(85) Each adjustment instruction or signal may comprise, but is not limited to, an indication of

whether the instruction is an increment instruction or a decrement instruction, the counter number, the task ID and, optionally the value of the increment or decrement. In some cases, some counters **103** may only be incremented or decremented by a fixed amount thus the value of the increment or decrement need not be explicitly provided. The counter update logic **404** uses the counter number and the task ID to identify a particular counter **103** in the matrix **402**. The counter update logic **404** then increments or decrements the value of the identified counter **103** based on the adjustment instructions (e.g. it increments or decrements the value of the counter by a specified or predetermined amount).

(86) Although the counter update logic **404** is shown as a single block that performs increments and decrements of the counters **103** in the matrix **402**, in other examples there may be separate increment and decrement logic blocks which receive and process increment and decrement instructions respectively. In these cases, the adjustment instructions may not specify whether the adjustment is an increment or a decrement since this can be automatically determined based on which logic block receives the adjustment instruction.

(87) The status logic **406** is configured to generate counter status information indicating the status of the counters **103** in the matrix **402** and provide the counter status information to at least the queues **110**. This allows the queues **110** to determine whether to stall a secondary instruction because one or more of the hazards related to a primary instruction from which it depends has not been resolved, or whether all hazards related to primary instructions from which it depends have been resolved and it is safe to forward, or issue, the secondary instruction to the instruction pipeline for execution (assuming there are no other hazards). Specifically, a secondary instruction can only be forwarded to, or issued to, an instruction pipeline **106** if the counters associated with the primary instruction(s) from which it depends indicate that the hazards related to the primary instructions from which it depends have been resolved.

(88) In some cases, the counter status information may comprise for example, a bit or flag for each counter that indicates, based on the value of the counter, whether there is a hazard related to the associated primary instruction or whether the hazard related to the associated primary instruction has been resolved. For example, a flag value “1” may indicate that there is a hazard related to the primary instruction and a flag value of “0” may indicate that the hazard related to the associated primary instruction has been resolved. The flag values may be generated based on the values of the counters **103**. For example, as described above, the counters **103** may be configured so that when they have one of one or more predetermined values (e.g. a non-zero value) they indicate that there is a hazard related to the associated primary instruction, and when they have one of one or more other predetermined values (e.g. a zero value) they indicate that the hazard related to the associated primary instruction has been resolved. In other cases, the counter status information may comprise the actual values of the counters **103** in the matrix **402** and the receiver of the counter status information (e.g. the queues **110**) is configured to determine from the counter values whether there is a hazard related to the associated primary instruction.

(89) In some cases, the status logic **406** may be configured to periodically push the counter status information to the queues **110**. In other cases, the queues **110** may be configured to request the counter status information and the status logic **406** may be configured to provide the counter status information in response to such a request.

(90) Reference is now made to FIG. 5 which illustrates an example method **500** for avoiding inter-pipeline data hazards in a parallel processing unit such as the parallel processing unit of FIG. 1. The method **500** begins at block **502** where the instruction decoder **104** receives a scheduled instruction to be decoded and at block **504** the instruction decoder **104** decodes the received instruction. If the instruction decoder **104** determines (at block **506**) that that decoded instruction is a primary instruction (e.g. from the primary instruction field) then the method **500** proceeds to block **508** where the instruction decoder **104** sends or transmits an adjustment instruction to the counter block **102** which causes the value of the counter **103** associated with the detected primary

instruction to be adjusted to have a value (e.g. a non-zero value) that indicates that there is a hazard related to the primary instruction. For example, the instruction decoder **104** may be configured to send or transmit an adjustment instruction to the counter block **102** that causes the value of the counter **103** associated with the detected primary instruction to be incremented by a predetermined amount (e.g. **8**). The method **500** then proceeds to block **510**.

(91) If the instruction decoder **104** determines at block **506** that the decoded instruction is not a primary instruction (e.g. from the primary instruction field) then the method **500** proceeds to block **510** where the instruction is forwarded to the queue **110** associated with the appropriate instruction pipeline **106**. Then at block **512** the queue **110** determines whether the received instruction is a secondary instruction. If the queue **110** determines (at block **512**) that the received instruction is a secondary instruction then the method **500** proceeds to block **514** where the instruction is stalled if at least one of the counter(s) associated with the primary instruction(s) from which the secondary instruction depends (e.g. as determined from the secondary instruction field) indicate that there is a hazard related to the primary instruction. Once the counter(s) associated with the primary instruction(s) from which the secondary instruction depends indicate that the hazard related to the primary instruction has been resolved the method **500** proceeds to block **516**.

(92) If the queue **110** determines (at block **512**) that the received instruction is not a secondary instruction then the method **500** proceeds directly to block **516** where the instruction is forwarded to the corresponding instruction pipeline for execution. The method **500** then proceeds to block **518**.

(93) At block **518** the monitor logic determines whether the instruction in the instruction pipeline is a primary instruction. If the monitor logic detects (at block **518**) that the instruction in the instruction pipeline is a primary instruction then the method **500** proceeds to block **520** where the monitor logic determines whether the instruction pipeline has resolved the hazard related to the primary instruction. If the monitor logic **108** determines that the instruction pipeline has resolved the hazard related to a primary instruction the monitor logic **108** sends (at block **522**) an adjustment signal or instruction to the counter block that causes the value of the counter associated with the primary instruction to be adjusted to have a value (e.g. zero) that indicates that hazard related to the primary instruction has been resolved.

(94) Where the number of counters per task ID is small compared to the number of inter-pipeline data hazards per task the counters may be frequently reused within a task. This may result in situations where two primary instructions associated with the same counter are received at the instruction decoder in relatively quick succession so that the instruction decoder may, when it forwards the second primary instruction for execution adjust the value of the common counter to indicate that the second primary instruction has a hazard before all of the secondary instructions related to the first primary instruction have had a chance to read the counter to see that the hazard related to the first primary instruction has been resolved. For example, if first and second primary instructions are associated with counter 1 and there are three secondary instructions that are dependent on the first primary instruction, it may be possible for the value of counter 1 to be adjusted to indicate that there is a hazard related to the first primary instruction in response to the instruction decoder forwarding the first primary instruction for execution, subsequently adjusted to indicate that the hazard related to the first primary instruction has been resolved in response to detecting the instruction pipeline has resolved the hazard, and subsequently adjusted to indicate that there is a hazard related to the second primary instruction in response to the instruction decoder forwarding the second primary instruction for execution, all before the queues of the three secondary instructions have read or obtained the status of counter 1. This would result in the secondary instructions having to wait for the second primary instruction to be completed. Having an instruction waiting on the effects of an instruction that occurs later in the code is clearly problematic as it could result in a deadlock.

(95) Accordingly, in some embodiments, to ensure that a counter **103** is not adjusted to reflect the status of a hazard related to a subsequent primary instruction until there are no secondary

instructions waiting on the result of the hazard related to a previous primary instruction, the counter block **102** further comprises a wait counter for each counter **103** that keeps track of the number of secondary instructions waiting on the results of the corresponding counter **103**. The value of the wait counter for a counter **103** may be adjusted (e.g. incremented) by the instruction decoder each time a secondary instruction that is dependent on a primary instruction associated with that counter is sent to a queue for execution to indicate that there is one additional secondary instruction waiting on the results of the counter **103**. When a secondary instruction that is dependent on a primary instruction associated with a counter **103** is sent from a queue to the corresponding instruction pipeline **106** for execution, or when a secondary instruction that is dependent on primary instruction associated with a counter **103** is discarded by a queue (e.g. in response to the queue determining the instruction is not active), the value of the wait counter for that counter may be adjusted (e.g. decremented) by the queue to indicate that there is one less secondary instruction waiting on the results of the counter **103**.

(96) When a primary instruction is received at the instruction decoder the instruction decoder determines if the wait counter for the corresponding counter **103** indicates that there are no secondary instructions waiting on the results of this counter. If the wait counter indicates that there are no secondary instructions waiting on the results of this counter then the instruction is forwarded to the appropriate queue for execution. If, however, the wait counter indicates that there is at least one secondary instruction waiting on the results of this counter then the instruction decoder causes that instruction and related instructions (e.g. instructions in the same task having the same task ID) to be de-scheduled (e.g. the instruction is sent back to a scheduler) where it waits until the wait counter indicates that there are no secondary instructions waiting on the results of the relevant counter.

(97) Reference is now made to FIG. **6** which illustrates a second example parallel processing unit **600** which may be a GPU or other parallel processing unit that comprises a plurality of wait counters. It will be appreciated that FIG. **6** only shows some elements of the parallel processing unit and there may be many other elements (e.g. caches, interfaces, etc.) within the parallel processing unit that are not shown in FIG. **6**. Like the parallel processing unit **100** of FIG. **1**, the parallel processing unit **600** of FIG. **6** comprises a counter block **602**, an instruction decoder **604**, a plurality of instruction pipelines **106**, monitor logic **108** and a queue **610** for each instruction pipeline **106**.

(98) The counter block **602** of FIG. **6** is that same as the counter block **102** of FIG. **1** except that the counter block **602** of FIG. **6** also comprises a wait counter **630** for each counter **103** to keep track of the number of secondary instructions waiting on the results of that counter **103**. In some examples, the wait counters **630** are configured so that a zero value indicates that there are no secondary instructions waiting on the results of the counter **103**, and a non-zero value indicates the number of secondary instructions waiting on the results of the counter **103** (e.g. a wait counter value of three indicates that there are three secondary instructions waiting on the results of the counter **103**). However, this is an example only and other values may be used to indicate the number of secondary instructions waiting on this counter.

(99) The counter block **602** is configured to adjust the values of the wait counters **630** in response to receiving wait counter adjustment instructions or signals from the instruction decoder **104**, and the queues **110**; and to generate and provide wait counter status information to the instruction decoder **604** and the scheduler **616**. In particular, as described in more detail below, the instruction decoder **604** is configured to send a wait counter adjustment instruction or signal to the counter block **602** that causes the counter block **602** to adjust the value of the wait counter **630** corresponding to a particular counter **103** to indicate that an additional secondary instruction is waiting on the results of the particular counter **103**, in response to forwarding a secondary instruction that is dependent on the primary instruction associated with the particular counter **103** to a queue **610**. The queues **610** are configured to send a wait counter adjustment instruction or

signal to the counter block **602** that causes the counter block **602** to adjust the value of the wait counter **630** corresponding to a particular counter **103** to indicate that one less secondary instruction is waiting on the results of the particular counter **103**, in response to (i) forwarding a secondary instruction that is dependent on the primary instruction associated with the particular counter **103** to an instruction pipeline for execution; or (ii) discarding a secondary instruction that is dependent on the primary instruction associated with the particular counter **103**.

(100) The wait counter status information comprises information that indicates whether there are any secondary instructions waiting on the results of a counter and is generated by the counter block **602** based on the value of the wait counters **630**. In some cases, the wait counter status information may comprise a flag or bit for each wait counter indicating whether there are any secondary instructions waiting on the results of the associated counter **103**. For example, the wait counter status information may comprise a flag for each wait counter where a flag is set to “1” to indicate that there is at least one secondary instruction waiting on the results of the associated counter and a flag is set to “0” to indicate that there are no secondary instructions waiting on the results of the associated counter. In other cases, the wait counter status information may comprise the value of each of the wait counters **630** and the recipient of the wait counter status information is configured to determine from the values whether there are any secondary instructions waiting on the results of the corresponding counter **103**.

(101) An example implementation of the counter block **602** of FIG. 6 is described below with reference to FIG. 7.

(102) The instruction decoder **604** of FIG. 6 is the same as the instruction decoder **104** of FIG. 1 except the instruction decoder **604** is further configured to, in response to forwarding a secondary instruction to a queue **610** for execution, send, for each primary instruction the secondary instruction is dependent on, a wait counter adjustment instruction to the counter block **602** that causes the counter block **602** to adjust the value of a particular wait counter to indicate that an additional secondary instruction is waiting on the results of the associated counter **103**. The instruction decoder **604** of FIG. 6 is also further configured to, prior to sending a decoded primary instruction to a queue **610** for execution, determine (from the wait queue status information) whether there are any secondary instructions waiting on the results of the counter associated with the primary instruction. If the instruction decoder **604** determines that there are no secondary instructions waiting on the results of the counter, the instruction decoder **604** sends the primary instruction to a queue **610** for execution. If, however, the instruction decoder **604** determines that there is at least one secondary instruction waiting on the results of the counter, the instruction decoder **604** sends a deactivate instruction or signal (which may also be referred as a de-schedule instruction or signal) to the scheduler **616** to cause the instruction (and/or task it is associated with) to be de-scheduled (or de-activated) until the wait counter status information indicates that there are no secondary instructions waiting on the results of the relevant counter. In some cases, the wait counter status information is automatically pushed to the instruction decoder **604** by the counter block **602** on a periodic basis. In other cases, the instruction decoder **604** is configured to request the wait counter status information from the counter block and in response to receiving the request the counter block **602** provides the wait counter status information to the instruction decoder **604**.

(103) The plurality of instruction pipelines **106** are the same as the instruction pipelines **106** of FIG. 1 and as such are configured to execute received instructions.

(104) The monitor logic **108** is the same as the monitor logic **108** of FIG. 1.

(105) The queues **610** of FIG. 6 are the same as the queues **110** of FIG. 1 except that each queue **610** of FIG. 6 is further configured to determine whether the instructions forwarded by the queue **610** to the corresponding instruction pipeline, or instructions discarded by the queue **610** (e.g. because they are inactive), are secondary instructions, and in response to the queue **610** forwarding a secondary instruction to the corresponding instruction pipeline or in response to the queue **610** discarding a secondary instruction, send, for each primary instruction that the secondary instruction

is dependent on, a wait counter adjustment instruction or signal to the counter block **602** that causes the counter block **602** to adjust the value of the wait counter corresponding to the counter associated with the primary instruction to indicate that one less secondary instruction is waiting on the results of that counter. In some cases, the wait counter adjustment instruction or signal sent by the queues **610** to the counter block **602** causes the counter block to decrement the value of the wait counter by one.

(106) The parallel processing unit **600** of FIG. **6** also includes a scheduler **616**. The scheduler **616** is the same as the scheduler **116** of FIG. **1** except that the instruction scheduler **616** is further configured to, in response to receiving a deactivate instruction or signal from the instruction decoder **604** for an instruction, to de-schedule or deactivate the identified instruction (and the associated task) until the relevant wait counter(s) indicate that there are no more secondary instructions waiting on the results of the associated counter.

(107) As described above, when a task has been sent from the scheduler **616** to the instruction fetch module **120** for execution that task becomes an active task. When the scheduler **616** receives a de-schedule or deactivate message for an instruction related to a task that task becomes an in-active task and is not available to become an active task until the relevant wait counter(s) indicates that there are no more secondary instructions waiting on the result of the corresponding counter. A deactivate or de-schedule instruction may comprise a task ID to identify the task to be de-scheduled/deactivated, a program counter value to indicate the instruction (so that the parallel processing unit can return to that instruction in the program when it is time for the task to be reactivated or rescheduled), and information identifying the relevant wait counter(s) to be monitored (e.g. a wait counter number).

(108) Reference is now made to FIG. **7** which illustrates an example implementation of the counter block **602** of FIG. **6**. The counter block **602** of FIG. **6** comprises the matrix **402** of counters **103**, counter update logic **404** and status logic **406** of FIG. **4** which operate as described above with reference to FIG. **4**. The counter block **602** of FIG. **6** additionally comprises a matrix **702** of wait counters **630**, wait counter update logic **704**, and wait counter status logic **706**.

(109) The matrix **702** of wait counters **630** comprises one wait counter **630** for each counter **103** in the matrix **402** of counters **103** which is used to track the number of secondary instructions that are waiting for the results of the corresponding counter **103**. The matrix **702** of wait counters **630** is arranged in the same manner as the matrix **402** so that the same information that is used to identify a counter **103** (e.g. task ID and counter number) can be used to identify the corresponding wait counter **630**. Specifically, the matrix **702** of wait counters comprises a row for each possible task ID. For example, where there are 3 possible task IDs the matrix **702** will have three rows as shown in FIG. **7**. Each row comprises K counters where K is the number of counters in each row of the matrix **402** of counters **103**.

(110) The wait counter update logic **704** receives wait counter adjustment instructions or signals from the instruction decoder **604** and the queues **610** which causes the value of a wait counter **630** to be adjusted to reflect the number of secondary instructions that are currently waiting on the results of the corresponding counter **103**. In particular, as described above in relation to FIG. **6**, in response to sending a secondary instruction to a queue **610** for execution the instruction decoder **604**, for each primary instruction the secondary instruction is dependent on, sends a wait counter adjustment instruction or signal to the counter block **602** to cause the value of a wait counter corresponding to the counter associated with the primary instruction to be adjusted to have a value that indicates that an additional secondary instruction is waiting on the results of the counter **103**. In some examples, the instruction decoder **604** is configured to transmit an adjustment instruction that causes the counter block **602** to increment the value of the wait counter **630** corresponding to the counter **103** by one.

(111) When a queue **610** subsequently detects that a secondary instruction has been sent to an instruction pipeline **106** for execution the queue **610** sends, for each primary instruction the

secondary instruction is dependent on, a wait counter adjustment instruction or signal to the counter block **602** to cause the value of the wait counter corresponding to the counter associated with the primary instruction to be adjusted to indicate that one less secondary instruction is waiting on the results of that counter. In some examples, each queue **610** is configured to transmit a wait counter adjustment instruction or signal to the counter block **602** that causes the counter block **602** to decrement the value of the wait counter **630** corresponding to the counter **103** by one.

(112) Also, if a wait queue **610** subsequently detects that a secondary instruction has been discarded (e.g. because it is inactive), the queue **610** sends, for each primary instruction the secondary instruction is dependent on, a wait counter adjustment instruction or signal to the counter block **602** to cause the value of the wait counter corresponding to the counter associated with the primary instruction to be adjusted to indicate that one less secondary instruction is waiting on the results of that counter.

(113) Each wait counter adjustment instruction or signal may comprise, but is not limited to, an indication of whether the instruction is an increment instruction or a decrement instruction, the counter number, the task ID and, optionally the value of the increment or decrement. In some cases, one or more wait counters **630** may only be incremented or decremented by a fixed amount thus the amount of the increment or decrement need not be explicitly provided. The wait counter update logic **704** uses the counter number and the task ID to identify a wait counter **630** in the matrix **702**. The wait counter update logic **704** then increments or decrements the value of the identified wait counter **630** based on the wait counter adjustment instructions (e.g. it increments or decrements the value of the counter by a specified or predetermined amount).

(114) Although the wait counter update logic **704** is shown as a single block that increments and decrements the wait counters **630** in the matrix **702**, in other examples there may be separate increment and decrement logic blocks which receive and process wait counter increment and decrement instructions respectively. In these cases, the wait counter adjustment instructions may not specify whether the adjustment is an increment or a decrement since this can be automatically determined based on which logic block receives the adjustment instruction.

(115) The wait counter status logic **706** is configured to generate wait counter status information from the value of the wait counters **630** that indicates the status of the wait counters **630** in the matrix **702** and provide the wait counter status information to at least the scheduler **616** and the instruction decoder **604**. Providing the wait counter status information to the instruction decoder **604** allows the instruction decoder **604** to determine whether it is safe for a primary instruction to be forwarded to a queue for execution. Providing the wait counter status information to the scheduler **616** allows the scheduler **616** to determine when the instruction/task can be re-scheduled.

(116) As described above, in some cases, the wait counter status information may comprise a bit or flag for each wait counter that indicates whether there are any secondary instructions waiting on the results of the corresponding counter **103**. For example, a flag value of “1” may indicate that there is at least one secondary instruction waiting on the results of the corresponding counter **103** and a flag value of “0” may indicate that there are no secondary instructions waiting on the results of the corresponding counter **103**. In other cases, the counter status information may comprise the actual values of the wait counters **630** in the matrix **702** and the receivers of the wait counter status information (e.g. the scheduler **616** and the instruction decoder **604**) are configured to determine from the counter values whether there are any secondary instructions that are waiting on the results of counter **103**.

(117) In some cases, the wait counter status logic **706** may be configured to periodically push the wait counter status information to the scheduler **616** and the instruction decoder **604**. In other cases, the scheduler **616** and the instruction decoder **604** may be configured to request the wait counter status information and the wait counter status logic **706** may be configured to provide the wait counter status information in response to such a request.

(118) Reference is now made to FIG. **8** which illustrates a second example method **800** for

avoiding inter-pipeline data hazards in a parallel processing unit, such as the parallel processing unit of FIG. 6. The method **800** begins at block **802** where an instruction decoder **604** receives an instruction and at block **804** decodes the instruction. At block **806** the instruction decoder **604** determines (e.g. from the primary instruction field) whether the decoded instruction is a primary instruction. If it is determined (at block **806**) that the instruction is a primary instruction then the method **800** proceeds to block **808** where a determination is made, based on the wait counter corresponding to the counter associated with the primary instruction, whether there are any secondary instructions waiting on the results of that counter. If it is determined (at block **808**) that there is at least one secondary instruction waiting on the results of the relevant counter then the method **800** proceeds to block **810** where the instruction decoder sends a de-schedule or deactivate instruction or signal to the scheduler **616** to cause the scheduler to deactivate the instruction/task until at block **812** the wait counter corresponding to the counter associated with the primary instruction indicates that there are no secondary instructions waiting on the results of the relevant counter. Once there are no secondary instructions waiting on the results of the relevant counter then at block **814** the instruction/task is re-scheduled and then the method **800** proceeds back to block **802**.

(119) If it is determined (at block **808**) that there are no secondary instructions waiting on the results of the relevant counter then it is safe to adjust the value of the counter for the new primary instruction so the method **800** proceeds to block **816** where the instruction decoder **604** sends or transmits an adjustment instruction to the counter block that causes the value of the counter associated with the primary instruction to be adjusted to indicate that there is a hazard related to the primary instruction. The method **800** then proceeds to block **818**.

(120) If it is determined (at block **806**) that the decoded instruction is not a primary instruction then the method proceeds to block **818** where the instruction decoder determines (e.g. from the secondary instruction field) whether the instruction is a secondary instruction. If it is determined that the instruction is a secondary instruction then at block **820** the instruction decoder **604** sends or transmits a wait counter adjustment instruction or signal to the counter block, for each primary instruction the secondary instruction is dependent thereon, that causes the value of the wait counter corresponding to the counter associated with the primary instruction to be adjusted to indicate that there is one more secondary instruction waiting on the results of that counter. The method **800** then proceeds to block **822**.

(121) If the instruction decoder **604** determines (at block **818**) that the decoded instruction is not a secondary instruction (e.g. from the secondary instruction field) then the method **800** proceeds to block **822** where the instruction is forwarded to the queue associated with the appropriate instruction pipeline. The method **800** then proceeds to block **824**.

(122) At block **824**, the queue determines whether the received instruction is a secondary instruction. If the queue determines (at block **824**) that the received instruction is a secondary instruction then the method **800** proceeds to block **826** where the instruction is stalled if at least one of the counter(s) associated with the primary instruction(s) from which the secondary instruction depends (e.g. as determined from the secondary instruction field) indicate that there is a hazard related to the primary instruction (e.g. the counter is non-zero). Once the counter(s) associated with the primary instruction(s) from which the secondary instruction depends indicate that the hazard related to the primary instruction has been resolved the method **800** proceeds to block **828** where the queue **610** transmits or sends a wait counter adjustment signal to the counter block, for each primary instruction the secondary instruction is dependent thereon, that causes the value of the wait counter corresponding to the counter associated with the primary instruction to be adjusted to indicate that there is one less secondary instruction waiting on the results of that counter. The method **800** then proceeds to block **830**.

(123) If the queue determines (at block **824**) that the received instruction is not a secondary instruction then the method **800** proceeds directly to block **830** where the instruction is forwarded

to the corresponding instruction pipeline **106** for execution. The method **800** then proceeds to block **832**.

(124) At block **832**, the monitor logic determines whether the instruction in the instruction pipeline is a primary instruction. If the monitor logic detects (at block **832**) that the instruction in the instruction pipeline is a primary instruction then the method **800** proceeds to block **834** where the monitor logic determines whether the instruction pipeline has resolved the hazard related to the primary instruction. If the monitor logic **108** determines that the instruction pipeline has resolved the hazard related to the primary instruction the monitor logic **108** sends (at block **836**) an adjustment signal or instruction to the counter block that causes the value of the counter associated with the primary instruction to be adjusted to have a value (e.g. zero) that indicates that hazard related to the primary instruction has been resolved.

(125) Stalling secondary instructions right before they are to be executed by an instruction pipeline has shown to improve performance in cases where the primary instruction(s) on which the secondary instruction depends will be completed quickly (e.g. when the primary instruction(s) are executed by an instruction pipeline with a high throughput). Such inter-pipeline data hazards are referred to herein as low latency inter-pipeline hazards. However, stalling secondary instructions right before they are to be executed by an instruction pipeline has shown to reduce performance where the primary instruction(s) on which the secondary instruction depends will be completed slowly (e.g. when the primary instruction(s) are executed by an instruction pipeline with low throughput). Such inter-pipeline data hazards may be referred to herein as high latency inter-pipeline hazards.

(126) Accordingly, in some embodiments described herein the compiler may be configured to separately identify and mark low latency inter-pipeline data hazards and high latency inter-pipeline data hazards. In these embodiments, low the low latency data hazards are processed as described above with reference FIG. **1** (i.e. secondary instructions that are waiting for one or more low latency hazards to clear will wait in a queue preceding the relevant instruction pipeline for the low latency hazard(s) to clear). However high latency data hazards are processed in a different manner. Specifically, if the instruction decoder detects a secondary instruction that is waiting for one or more high latency hazards to clear, instead of sending the secondary instruction to a queue for execution, the instruction is de-scheduled until all relevant high latency inter-pipeline hazards have cleared.

(127) Reference is now made to FIG. **9** which illustrates a third example parallel processing unit **900** which may be a GPU or other parallel processing unit that processes high and low latency inter-pipeline data hazards in different manners. It will be appreciated that FIG. **9** only shows some elements of the parallel processing unit and there may be many other elements (e.g. caches, interfaces, etc.) within the parallel processing unit **900** that are not shown in FIG. **9**. Like the parallel processing unit **100** of FIG. **1**, the parallel processing unit **900** of FIG. **9** comprises a counter block **102**, an instruction decoder **904**, a plurality of instruction pipelines **106**, monitor logic **108**, and a queue **910** for each instruction pipeline **106**.

(128) The parallel processing unit **900** of FIG. **9** is configured to process instructions that have been generated by a compiler that is configured to (i) identify inter-pipeline data hazards; (ii) determine whether the identified inter-pipeline data hazard is a low latency data hazard or a high latency data hazard (e.g. based on the throughput associated with the instruction pipeline that will execute the primary instruction of the inter-pipeline data hazard); and (iii) allocate a counter **103** from a first subset of the counters **103** (e.g. the low latency counters) to a low latency inter-pipeline data hazard and allocate a counter **103** from a second, different, subset of the counters **103** (e.g. the high latency counters) to a high latency inter-pipeline data hazard. Accordingly, the instructions that are processed by the parallel processing unit **900** of FIG. **9** include information (inserted at build time—e.g. by a compiler) that identify primary instructions (i.e. instructions from which at least one other instruction in another instruction pipeline is dependent on), secondary instructions

(i.e. instructions that are dependent on at least one primary instruction in another pipeline), the counter(s) they are associated with, and information indicating whether or not the associated hazard is a low latency hazard or a high latency hazard.

(129) For example, as shown in FIG. 10, if there are three counters numbered 1 to 3 (e.g. for each possible task ID) then a first set of counters formed by counters 1 and 2 may be allocated to low latency inter-pipeline data hazards and a second set of counters formed by counter 3 may be allocated to high latency inter-pipeline data hazards. The counter number associated with a primary instruction or a secondary instruction can be used by the parallel processing unit 900 to determine whether the primary instruction or the secondary instruction is associated with a low latency inter-pipeline data hazard or a high latency inter-pipeline data hazard and process it accordingly.

(130) Returning to FIG. 9, the counter block 102 is the same as the counter block 102 of FIG. 1 except that the counter block 102 is further configured to provide counter status information to the scheduler 916 and the instruction decoder 904. Specifically, the counter block 102 comprises a plurality of counters 103 which are used to track inter-pipeline data hazards. The counter block 102 receives an adjustment instruction or signal from the instruction decoder 904 to adjust the value of a counter to indicate that there is a hazard related to the associated primary instruction when the instruction decoder forwards a primary instruction to a queue 910 for execution; and the counter block 102 subsequently receives an adjustment instruction or signal from the monitor logic 108 to adjust the value of the counter to indicate that the hazard related to the associated primary instruction has been resolved when the monitor logic 108 detects that an instruction pipeline has resolved a hazard related to a primary instruction. The counter block also provides counter status information to the queues 910, the scheduler 916 and the instruction decoder 904 that indicates, for each counter, whether there is hazard related to the primary instruction associated therewith. The fact that one or more of the counters 103 is/are designated for use with low latency inter-pipeline data hazards and one or more of the counters 103 is/are designated for use with high latency inter-pipeline data hazards does not affect the operation of the counter block 102.

(131) The instruction decoder 904 generally corresponds to the instruction decoder 104 of FIG. 1 in that it is configured to receive instructions, decode the received instructions, forward the decoded instructions to a queue 910 associated with an appropriate pipeline for execution; and in response to forwarding a primary instruction (whether related to a low latency inter-pipeline data hazard or a high latency inter-pipeline data hazard) to a queue 910 for execution send an adjustment instruction to the counter block to cause the value of the counter associated with the primary instruction to be adjusted to indicate that there is hazard related to the primary instruction (e.g. send an adjustment instruction to increment the value of the counter associated with the primary instruction by a predetermined amount (e.g. 8)).

(132) However, instead of automatically forwarding all decoded secondary instructions to queues 910 for execution like the instruction decoder 104 of FIG. 1, the instruction decoder 904 of FIG. 9 is configured to, prior to forwarding a decoded secondary instruction to a queue 910 for execution, determine (e.g. from the counter numbers associated therewith) whether the secondary instruction is related to at least one high latency data hazard. If it is determined that the secondary instruction is not related to any high latency data hazards the instruction decoder 904 forwards the secondary instruction to a queue 910 for execution.

(133) If, however, it is determined that the secondary instruction is related to at least one high latency data hazard then the instruction decoder 904 determines from the counter status information provided by the counter block 102 whether each high latency inter-pipeline data hazard associated with the secondary instruction has been resolved. If the counter status information indicates that all high latency data hazards associated with the secondary instruction have been resolved (e.g. any high latency counters associated with the secondary instruction indicate that the high latency data hazard has been resolved) then the secondary instruction is forwarded to a queue 910 for execution. If, however, the counter status information indicates that at least one of the high latency hazards

associated with the secondary instruction has not been resolved then a deactivation instruction or signal is sent to the scheduler **916** to cause the scheduler to deactivate the instruction (and/or task associated therewith). The deactivation instruction may comprise information identifying the instruction (e.g. the value of the program counter) and information identifying the high latency counters to be monitored, and where the instruction is part of a task, the deactivation message may also comprise information identifying the task (e.g. a task ID).

(134) The plurality of instruction pipelines **106** are the same as the instruction pipelines **106** of FIG. **1** and as such are configured to execute received instructions.

(135) The monitor logic **108** is the same as the monitor logic **108** of FIG. **1**.

(136) The queues **910** are the same as the queues **110** of FIG. **1** except that the queues **910** of FIG. **9** are only configured to stall a secondary instruction while a low latency counter (instead of any counter) associated with the secondary instruction indicates there is a hazard related to the associated primary instruction.

(137) The parallel processing unit **900** of FIG. **9** also includes a scheduler **916**. The scheduler **916** is the same as the scheduler **116** of FIG. **1** except that the scheduler **916** of FIG. **9** is further configured to, in response to receiving a deactivate instruction or signal from the instruction decoder **904** for a secondary instruction, de-schedule or deactivate the identified secondary instruction (and the associated task) until the high latency counters associated with the secondary instruction indicate that the high latency data hazard(s) has/have been resolved.

(138) As described above, when a task has been sent from the scheduler **916** to the instruction fetch module for processing that task becomes an active task. When the scheduler **916** receives a de-schedule or deactivate message for an instruction the associated task becomes an in-active task and is not available to become an active task until the relevant high latency counter(s) indicate that the relevant high latency hazards have been resolved. A deactivate or de-schedule instruction may comprise information identifying the instruction (e.g. a program counter value) to be de-scheduled/deactivated, information identifying the task the instruction is associated with (e.g. a task ID), and information identifying the high latency counters to be monitored (e.g. counter numbers).

(139) Reference is now made to FIG. **11** which illustrates a third example method **1100** for avoiding inter-pipeline data hazards in a parallel processing unit, such as the parallel processing unit **900** of FIG. **9**. The method **1100** begins at block **1102** where an instruction decoder **904** receives an instruction and decodes (at block **1104**) the instruction. At block **1106**, the instruction decoder **904** determines (e.g. from the secondary instruction field) whether the decoded instruction is a secondary instruction. If (at block **1106**) it is determined that the instruction is a secondary instruction then the method **1100** proceeds to block **1108** where a determination is made, from the number of the counter associated with each primary instruction from which it depends whether the instruction relates to at least one high latency inter-pipeline hazard. If it is determined (at block **1108**) that the secondary instruction relates to at least one high latency inter-pipeline data hazard then the method **1100** proceeds to block **1110** where the instruction decoder **904** determines if the relevant high latency counters from which the secondary instruction depends indicate that there is a high latency hazard related to the primary instruction. If (at block **1110**) it is determined that the relevant high latency counters indicate that there is a high latency hazard related to the primary instruction the method **1100** proceeds to block **1112** where the instruction decoder **904** sends a de-schedule or deactivate instruction or message to the scheduler which causes the scheduler to deactivate the instruction/task until at block **1114** the relevant high latency counter(s) indicate that the high latency hazard related to the primary instruction(s) has/have been resolved. Once the relevant high latency counters indicate that the high latency hazard related to the primary instruction(s) have been resolved then at block **1116** the instruction/task is re-scheduled and the method **1100** proceeds back to block **1102**.

(140) If it is determined (at block **1108**) that the secondary instruction is not associated with any high latency counters, or (at block **1110**) that the relevant high latency hazard(s) have been resolved

the method proceeds to block **1118**.

(141) If it is determined (at block **1106**) that the decoded instruction is not a secondary instruction then the instruction decoder determines (at block **1118**) (e.g. from the primary instruction field) whether the instruction is a primary instruction. If it is determined that the instruction is a primary instruction then at block **1120** the instruction decoder **904** sends or transmits a counter adjustment instruction or signal to the counter block that causes the value of the counter associated with the primary instruction to be adjusted to indicate that there is a hazard related to the primary instruction. The method (**1100**) then proceeds to block **1122**.

(142) If the instruction decoder determines (at block **1118**) that the decoded instruction is not a primary instruction (e.g. from the primary instruction field) then the method **1100** proceeds to block **1122** where the instruction is forwarded to the queue associated with the appropriate instruction pipeline. Then at block **1124** the queue **910** determines whether the received instruction is a secondary instruction. If the queue **910** determines (at block **1124**) that the received instruction is a secondary instruction then the method **1100** proceeds to block **1126** where the instruction is stalled until the low latency counter(s) associated with the primary instruction(s) from which the secondary instruction depends indicate that the low latency hazard related to the primary instruction has been resolved (e.g. the counter is zero). Once the low latency counter(s) associated with the primary instruction(s) from which the secondary instruction depends indicate that the low latency hazard related to the primary instruction has been resolved the method **1100** proceeds to block **1128**.

(143) If the queue **910** determines (at block **1124**) that the received instruction is not a secondary instruction then the method **1100** proceeds directly to block **1128** where the instruction is forwarded to the corresponding instruction pipeline for execution.

(144) At block **1130**, the monitor logic determines whether the instruction in the instruction pipeline is a primary instruction. If the monitor logic determines (at block **1130**) that the instruction in the instruction pipeline is a primary instruction then the method **1100** proceeds to block **1132** where the monitor logic determines whether the instruction pipeline has resolved the hazard related to the primary instruction. If the monitor logic **108** determines that the instruction pipeline has resolved the hazard related to the primary instruction the monitor logic **108** sends (at block **1134**) an adjustment signal or instruction to the counter block that causes the value of the counter associated with the primary instruction to be adjusted to indicate that the hazard related to the primary instruction has been resolved.

(145) Although not shown, in other examples the wait counters described with reference to FIGS. **6** to **8** may be combined with the low latency/high latency processing of inter-pipeline data hazards described with reference FIGS. **9** to **11**. In these examples, there may only be a wait counter for each of the low latency counters.

(146) Reference is now made to FIG. **12** which illustrates an example method **1200** for generating instructions to be executed by the parallel processing unit **100**, **600** of FIGS. **1** and **6**. The method **1200** begins at block **1202** where a set of related instructions (e.g. instructions forming a program) are received (e.g. at a compiler). The set of related instructions may be in any suitable format. For example, in some cases the set of related instructions may be in a machine language. Once the set of related instructions have been received the method **1200** proceeds to block **1204**.

(147) At block **1204**, the instructions are analysed (e.g. by the compiler) to identify inter-pipeline data hazards in the set of related instructions. As described above, an inter-pipeline data hazard exists when an instruction to be executed in one pipeline can only safely execute after another instruction to be executed in a different instruction pipeline has performed a data action (e.g. a read or write of data). There are generally three types of data hazards that may occur—a WAR hazard, a RAW hazard and WAW hazard.

(148) The instruction that performs the action is referred to herein as the primary instruction and the instruction that is waiting on the action by the primary instruction is referred to as the

secondary instruction. For example, in a WAR hazard the read instruction is the primary instruction and the write instruction is the secondary instruction; in a RAW hazard the write instruction is the primary instruction and the read instruction is the secondary instruction; and in a WAW hazard the first write instruction is the primary instruction and the second write instruction is the secondary instruction. The secondary instruction is said to be dependent on the primary instruction. As described above, there may be more than one secondary instruction that is dependent on the same primary instruction.

(149) In some cases, the compiler may be configured to identify inter-pipeline data hazards (and the primary and secondary instruction thereof) in the set of related instructions by analysing the operands of the instructions to identify WAR, RAW and WAW hazards. In other cases, the set of instructions may include information (e.g. manually inserted by a programmer or developer) indicating inter-pipeline data hazards in the set of instructions and the compiler may be configured to identify the inter-pipeline data hazards (and the primary and secondary instructions thereof) in the set of instructions based on this information. Once the inter-pipeline data hazards have been identified the method **1200** proceeds to block **1206**.

(150) At block **1206**, each primary instruction of an identified inter-pipeline data hazard is allocated (e.g. by the compiler) a counter (e.g. a counter number) of a plurality of counters (e.g. counters **103**) for tracking the identified inter-pipeline data hazard. As described above, in some cases the counters may be allocated to the primary instructions in a manner that maximizes the time between reuse of the counters (e.g. in a round-robin manner). Once a counter has been allocated to each primary instruction the method **1200** proceeds to block **1208**.

(151) At block **1208**, a computer executable instruction is generated (e.g. by the compiler) for each primary instruction such that the computer executable instruction comprises information that identifies the primary instruction as a primary instruction and information that identifies the counter allocated to the primary instruction. As described above with respect to FIG. 2, in some cases the generated computer executable instructions may comprise a primary instruction field that comprises the information indicating the computer executable instruction is a primary instruction and the information identifying the counter allocated to the primary instruction. For example, the primary instruction field may be configured to hold a number and when the number is a predetermined value (e.g. zero) it indicates that the instruction is not a primary instruction and when the number is not the predetermined value (e.g. a non-zero value) it indicates that the instruction is a primary instruction and the number represents a number of the counter associated with the primary instruction. Once computer readable instructions have been generated for the primary instructions, the method **1200** proceeds to block **1210**.

(152) At block **1210**, a computer executable instruction is generated (e.g. by the compiler) for each secondary instruction such that the computer executable instruction comprises information identifying the computer executable instruction as a secondary instruction and information identifying the counter allocated to the corresponding primary instruction. As described above with respect to FIG. 2, in some cases the generated computer executable instructions comprise a secondary instruction field that comprises the information indicating the computer executable instruction is a secondary instruction and the information identifying the counter allocated to the corresponding primary instruction. For example, the secondary instruction field may be configured to hold a bit mask wherein each bit of the bit mask corresponds to a counter of the plurality of counters and when a bit of the mask is set it indicates that the instruction is a secondary instruction that is dependent on the primary instruction associated with the corresponding counter. Once computer readable instructions have been generated for the secondary instructions, the method **1200** proceeds to block **1212**.

(153) At block **1212**, the computer executable instructions generated in blocks **1208** and **1210** are loaded into memory of, or accessible by, the parallel processing unit.

(154) Although in the example method **1200** of FIG. 12 all inter-pipeline data hazards are identified

and then counters are allocated thereto and executable instructions generated for the primary and secondary instructions thereof, in other examples, after each inter-pipeline data hazard has been identified a counter may be allocated to the primary instruction and the computer executable instructions are generated for the primary and secondary instructions to include information that indicates them as such and identifies the counter allocated to the primary instruction.

(155) When the instructions generated by the method **1200** of FIG. **12** are intended to be executed by a parallel processing unit, such as the parallel processing unit of FIG. **9**, which is configured to process low latency inter-pipeline data hazards differently than high latency inter-pipeline data hazards, the counters used for tracking the data hazards may be divided into two sets and counters from one set may be allocated to the primary instructions of low latency data hazards and counters from the other set may be allocated to the primary instructions of high latency data hazards. As described above, the parallel processing unit can then determine whether to process a hazard as a low latency hazard or a high latency hazard based on the counter allocated to the corresponding primary instruction.

(156) Reference is now made to FIG. **13** which illustrates an example method of allocating counters to the primary instructions (block **1206**) wherein primary instructions of high latency data hazards are allocated counters from one set of counters and primary instructions of low latency data hazards are allocated counters from another set of counters. The method **1206** begins at block **1302** where a determination is made for an identified inter-pipeline data hazard whether the inter-pipeline data hazard is a low latency data hazard or a high latency data hazard. The determination may be based on the throughput of the instruction pipeline (e.g. instruction pipeline **106**) that will execute the primary instruction of the identified inter-pipeline data hazard. The throughput provides an indicator of how quickly the instruction pipeline processes instructions. Generally, the higher the throughput the more quickly an instruction pipeline processes instructions. An instruction pipeline with a throughput above, or equal, to a predetermined threshold may be considered a high throughput instruction pipeline and a primary instruction that is to be executed by a high throughput instruction pipeline may be considered to be related to a low latency inter-pipeline data hazard. Conversely an instruction pipeline with a throughput below the predetermined threshold may be considered a low throughput instruction pipeline and a primary instruction that is to be executed by a low throughput instruction pipeline may be considered to be related to a high latency inter-pipeline data hazard. The determination of which instruction pipeline will execute the primary instruction may be based on the type of primary instruction, which may be identified, for example, by the operation code (“opcode”) of the primary instruction.

(157) If it is determined that the primary instruction is associated with a low latency hazard the method **1206** proceeds to block **1304** where the primary instruction is allocated (e.g. by the compiler) a counter from a first set of counters (e.g. counters 0 to 3). The method **1206** then proceeds to block **1308**. If, however, it is determined that the primary instruction is associated with a high latency hazard the method **1206** proceeds to block **1306** where the primary instruction is allocated (e.g. by the compiler) a counter from a second set of counters (e.g. counters 4 to 7). The method **1306** then proceeds to block **1308**.

(158) At block **1308** a determination is made as to whether there are any more identified inter-pipeline data hazards. If, there is at least one more identified inter-pipeline data hazard the method proceeds back to block **1302**. If there are no more identified inter-pipeline data hazards (i.e. all primary instructions have been allocated a counter) then the method **1206** ends.

(159) FIG. **14** shows a computer system in which the parallel processing units **100**, **600**, **900** described herein may be implemented. The computer system comprises a CPU **1402**, a GPU **1404**, a memory **1406** and other devices **1414**, such as a display **1416**, speakers **1418** and a camera **1420**. The parallel processing unit **1410** (corresponding to any of the parallel processing units **100**, **600** or **900**) is shown implemented within the GPU **1404**. In other examples, the parallel processing unit **1410** may be implemented within the CPU **1402**. The components of the computer system can

communicate with each other via a communications bus 1422.

(160) The parallel processing units, queues, and counter blocks of FIGS. 1, 3, 4, 6, 7 and 9 are shown as comprising a number of functional blocks. This is schematic only and is not intended to define a strict division between different logic elements of such entities. Each functional block may be provided in any suitable manner. It is to be understood that intermediate values described herein as being formed by an element of the parallel processing unit, queue, or counter blocks need not be physically generated by the parallel processing unit, queue or counter block at any point and may merely represent logical values which conveniently describe the processing performed by the parallel processing unit, queue, or counter block between its input and output.

(161) The parallel processing units, queues, and/or counter blocks described herein may be embodied in hardware on an integrated circuit. The parallel processing units described herein may be configured to perform any of the methods described herein. Generally, any of the functions, methods, techniques or components described above can be implemented in software, firmware, hardware (e.g., fixed logic circuitry), or any combination thereof. The terms “module,” “functionality,” “component,” “element,” “unit,” “block” and “logic” may be used herein to generally represent software, firmware, hardware, or any combination thereof. In the case of a software implementation, the module, functionality, component, element, unit, block or logic represents program code that performs the specified tasks when executed on a processor. The algorithms and methods described herein could be performed by one or more processors executing code that causes the processor(s) to perform the algorithms/methods. Examples of a computer-readable storage medium include a random-access memory (RAM), read-only memory (ROM), an optical disc, flash memory, hard disk memory, and other memory devices that may use magnetic, optical, and other techniques to store instructions or other data and that can be accessed by a machine.

(162) The terms computer program code and computer readable instructions as used herein refer to any kind of executable code for processors, including code expressed in a machine language, an interpreted language or a scripting language. Executable code includes binary code, machine code, bytecode, code defining an integrated circuit (such as a hardware description language or netlist), and code expressed in a programming language code such as C, Java or OpenCL. Executable code may be, for example, any kind of software, firmware, script, module or library which, when suitably executed, processed, interpreted, compiled, executed at a virtual machine or other software environment, cause a processor of the computer system at which the executable code is supported to perform the tasks specified by the code.

(163) A processor, parallel processing unit, computer, or computer system may be any kind of device, machine or dedicated circuit, or collection or portion thereof, with processing capability such that it can execute instructions. A processor may be any kind of general purpose or dedicated processor, such as a CPU, GPU, System-on-chip, state machine, media processor, an application-specific integrated circuit (ASIC), a programmable logic array, a field-programmable gate array (FPGA), or the like. A computer or computer system may comprise one or more processors.

(164) It is also intended to encompass software which defines a configuration of hardware as described herein, such as HDL (hardware description language) software, as is used for designing integrated circuits, or for configuring programmable chips, to carry out desired functions. That is, there may be provided a computer readable storage medium having encoded thereon computer readable program code in the form of an integrated circuit definition dataset that when processed in an integrated circuit manufacturing system configures the system to manufacture a system configured to perform any of the methods described herein, or to manufacture a processor or a parallel processing unit comprising any apparatus described herein. An integrated circuit definition dataset may be, for example, an integrated circuit description.

(165) There may be provided a method of manufacturing, at an integrated circuit manufacturing system, a parallel processing unit, queue and/or counter block as described herein. There may be

provided an integrated circuit definition dataset that, when processed in an integrated circuit manufacturing system, causes the method of manufacturing a parallel processing unit, queue and/or counter block to be performed.

(166) An integrated circuit definition dataset may be in the form of computer code, for example as a netlist, code for configuring a programmable chip, as a hardware description language defining an integrated circuit at any level, including as register transfer level (RTL) code, as high-level circuit representations such as Verilog or VHDL, and as low-level circuit representations such as OASIS® and GDSII. Higher level representations which logically define an integrated circuit (such as RTL) may be processed at a computer system configured for generating a manufacturing definition of an integrated circuit in the context of a software environment comprising definitions of circuit elements and rules for combining those elements in order to generate the manufacturing definition of an integrated circuit so defined by the representation. As is typically the case with software executing at a computer system so as to define a machine, one or more intermediate user steps (e.g. providing commands, variables etc.) may be required in order for a computer system configured for generating a manufacturing definition of an integrated circuit to execute code defining an integrated circuit so as to generate the manufacturing definition of that integrated circuit.

(167) An example of processing an integrated circuit definition dataset at an integrated circuit manufacturing system so as to configure the system to manufacture a parallel processing unit (e.g. parallel processing unit **100**, **600** or **900**) will now be described with respect to FIG. **15**.

(168) FIG. **15** shows an example of an integrated circuit (IC) manufacturing system **1502** which is configured to manufacture a parallel processing unit (e.g. parallel processing unit **100**, **600** or **900**) as described in any of the examples herein. In particular, the IC manufacturing system **1502** comprises a layout processing system **1504** and an integrated circuit generation system **1506**. The IC manufacturing system **1502** is configured to receive an IC definition dataset (e.g. defining a parallel processing unit (e.g. parallel processing unit **100**, **600** or **900**) as described in any of the examples herein), process the IC definition dataset, and generate an IC according to the IC definition dataset (e.g. which embodies a parallel processing unit (e.g. parallel processing unit **100**, **600** or **900**) as described in any of the examples herein). The processing of the IC definition dataset configures the IC manufacturing system **1502** to manufacture an integrated circuit embodying a parallel processing unit (e.g. parallel processing unit **100**, **600** or **900**) as described in any of the examples herein.

(169) The layout processing system **1504** is configured to receive and process the IC definition dataset to determine a circuit layout. Methods of determining a circuit layout from an IC definition dataset are known in the art, and for example may involve synthesising RTL code to determine a gate level representation of a circuit to be generated, e.g. in terms of logical components (e.g. NAND, NOR, AND, OR, MUX and FLIP-FLOP components). A circuit layout can be determined from the gate level representation of the circuit by determining positional information for the logical components. This may be done automatically or with user involvement in order to optimise the circuit layout. When the layout processing system **1504** has determined the circuit layout it may output a circuit layout definition to the IC generation system **1506**. A circuit layout definition may be, for example, a circuit layout description.

(170) The IC generation system **1506** generates an IC according to the circuit layout definition, as is known in the art. For example, the IC generation system **1506** may implement a semiconductor device fabrication process to generate the IC, which may involve a multiple-step sequence of photo lithographic and chemical processing steps during which electronic circuits are gradually created on a wafer made of semiconducting material. The circuit layout definition may be in the form of a mask which can be used in a lithographic process for generating an IC according to the circuit definition. Alternatively, the circuit layout definition provided to the IC generation system **1506** may be in the form of computer-readable code which the IC generation system **1506** can use to form a suitable mask for use in generating an IC.

(171) The different processes performed by the IC manufacturing system **1502** may be implemented all in one location, e.g. by one party. Alternatively, the IC manufacturing system **1502** may be a distributed system such that some of the processes may be performed at different locations, and may be performed by different parties. For example, some of the stages of: (i) synthesising RTL code representing the IC definition dataset to form a gate level representation of a circuit to be generated, (ii) generating a circuit layout based on the gate level representation, (iii) forming a mask in accordance with the circuit layout, and (iv) fabricating an integrated circuit using the mask, may be performed in different locations and/or by different parties.

(172) In other examples, processing of the integrated circuit definition dataset at an integrated circuit manufacturing system may configure the system to manufacture a parallel processing unit (e.g. parallel processing unit **100**, **600** or **900**) without the IC definition dataset being processed so as to determine a circuit layout. For instance, an integrated circuit definition dataset may define the configuration of a reconfigurable processor, such as an FPGA, and the processing of that dataset may configure an IC manufacturing system to generate a reconfigurable processor having that defined configuration (e.g. by loading configuration data to the FPGA).

(173) In some embodiments, an integrated circuit manufacturing definition dataset, when processed in an integrated circuit manufacturing system, may cause an integrated circuit manufacturing system to generate a device as described herein. For example, the configuration of an integrated circuit manufacturing system in the manner described above with respect to FIG. **15** by an integrated circuit manufacturing definition dataset may cause a device as described herein to be manufactured.

(174) In some examples, an integrated circuit definition dataset could include software which runs on hardware defined at the dataset or in combination with hardware defined at the dataset. In the example shown in FIG. **15**, the IC generation system may further be configured by an integrated circuit definition dataset to, on manufacturing an integrated circuit, load firmware onto that integrated circuit in accordance with program code defined at the integrated circuit definition dataset or otherwise provide program code with the integrated circuit for use with the integrated circuit.

(175) The implementation of concepts set forth in this application in devices, apparatus, modules, and/or systems (as well as in methods implemented herein) may give rise to performance improvements when compared with known implementations. The performance improvements may include one or more of increased computational performance, reduced latency, increased throughput, and/or reduced power consumption. During manufacture of such devices, apparatus, modules, and systems (e.g. in integrated circuits) performance improvements can be traded-off against the physical implementation, thereby improving the method of manufacture. For example, a performance improvement may be traded against layout area, thereby matching the performance of a known implementation but using less silicon. This may be done, for example, by reusing functional blocks in a serialised fashion or sharing functional blocks between elements of the devices, apparatus, modules and/or systems. Conversely, concepts set forth in this application that give rise to improvements in the physical implementation of the devices, apparatus, modules, and systems (such as reduced silicon area) may be traded for improved performance. This may be done, for example, by manufacturing multiple instances of a module within a predefined area budget.

(176) The applicant hereby discloses in isolation each individual feature described herein and any combination of two or more such features, to the extent that such features or combinations are capable of being carried out based on the present specification as a whole in the light of the common general knowledge of a person skilled in the art, irrespective of whether such features or combinations of features solve any problems disclosed herein. In view of the foregoing description it will be evident to a person skilled in the art that various modifications may be made within the scope of the invention.

(177) Further examples are now provided by way of clauses. Clause 1. A computer-implemented

method **1200** of generating computer executable instructions for a parallel processing unit, the method **1200** comprising, by a processor: receiving **1202** a plurality of related instructions; identifying **1204** data hazards in the plurality of related instructions, each data hazard comprising a primary instruction and one or more secondary instructions; allocating **1206** each primary instruction a counter of a plurality of counters for tracking the identified data hazard; generating **1208** a computer executable instruction for each primary instruction that comprises information indicating that the computer executable instruction is a primary instruction and information identifying the counter allocated to the primary instruction; and generating **1210** a computer executable instruction for each secondary instruction that comprises information indicating that the computer executable instruction is a secondary instruction and information identifying the counter allocated to the corresponding primary instruction; and loading **1212** the computer executable instructions into the parallel processing unit. Clause 2. The method **1200** of clause 1, wherein each computer executable instruction comprises a primary instruction field and a secondary instruction field, the primary instruction field configured to comprise the information indicating that the computer executable instruction is a primary instruction and the information identifying the counter allocated to the primary instruction, and the secondary instruction field configured to comprise the information indicating that the computer executable instruction is a secondary instruction and the information identifying the counter allocated to the corresponding primary instruction. Clause 3. The method **1200** of clause 2, wherein the primary instruction field is configured to hold a number and when the number is a predetermined value it indicates that the instruction is not a primary instruction and when the number is not the predetermined value it indicates that the instruction is a primary instruction and the number represents a number of the counter allocated to the primary instruction. Clause 4. The method **1200** of clause 2 or clause 3, wherein the secondary instruction field is configured to hold a bit mask wherein each bit of the bit mask corresponds to a counter of the plurality of counters and when a bit of the mask is set it indicates that the instruction is a secondary instruction that is dependent on the primary instruction allocated the corresponding counter. Clause 5. The method **1200** of any of clauses 2 to 4, wherein allocating a primary instruction a counter for tracking the identified data hazard comprises determining **1302** whether the identified data hazard is a high latency data hazard or a low latency data hazard, and allocating **1304** the primary instruction a counter from a first subset of the plurality of counters when the identified data hazard is a high latency data hazard, and allocating **1306** the primary instruction a counter from a second subset of the plurality of counters when the identified data hazard is a low latency data hazard. Clause 6. The method **1200** of clause 5, wherein the determination of whether an identified data hazard is a high latency data hazard or a low latency data hazard is based on a throughput associated with an instruction pipeline to execute the primary instruction of the identified data hazard. Clause 7. The method **1200** of clause 6, wherein an identified data hazard is a high latency data hazard if the throughput associated with the instruction pipeline to execute the primary instruction of the identified data hazard is greater than or equal to a predetermined threshold and the identified data hazard is a low latency data hazard otherwise. Clause 8. The method **1200** of clause 6 or clause 7, wherein the instruction pipeline to execute the primary instruction of the identified data hazard is based on a type of the primary instruction. Clause 9. The method **1200** of any of clauses 1 to 8, wherein the parallel processing unit comprises a plurality of parallel instruction pipelines, and identifying data hazards within the plurality of related instructions comprises identifying inter-pipeline data hazards within the plurality of related instructions. Clause 10. The method **1200** of any of clauses 1 to 9, wherein the parallel processing unit is configured to track data hazards using the counters identified in the computer readable instructions. Clause 11. The method **1200** of any of clauses 1 to 10, wherein the parallel processing unit is configured to: in response to a decoder of the parallel processing unit outputting a primary instruction for execution, adjust the counter allocated to the primary instruction to indicate a data hazard related to that primary instruction; in response to detecting that the data hazard related to a

primary instruction has been resolved, adjust the counter allocated to the primary instruction to indicate the data hazard has been resolved; and in response to the decoder outputting a secondary instruction for execution, stalling the secondary instruction if at least one counter allocated to a primary instruction from which the secondary instruction depends indicates there is data hazard related to the primary instruction. Clause 12. The method **1200** of any of clauses 1 to 11, wherein the parallel processing unit comprises: a plurality of counters **103**; a plurality of queues **110, 610, 910**, each queue **110, 610, 910** preceding one instruction pipeline of a plurality of instruction pipelines **106**; an instruction decoder **104, 604, 904** configured to: decode a received instruction; in response to determining the decoded instruction is a primary instruction from which at least one other instruction is dependent on, cause a value of a counter **103** of the plurality of counters allocated to the primary instruction to be adjusted to indicate that there is a hazard related to the primary instruction; and forward the decoded instruction to one of the plurality of queues **110, 610, 910**; and monitor logic **108** configured to monitor the plurality of instruction pipelines **106**, and in response to detecting that an instruction pipeline **106** has resolved a hazard related to a primary instruction, cause the value of the counter allocated to the primary instruction to be adjusted to indicate that the hazard related to the primary instruction has been resolved; wherein each queue **110, 610, 910** is configured to, in response to receiving a secondary instruction that is dependent on one or more primary instructions, stall execution of the secondary instruction by the associated instruction pipeline **106** if a counter **103** allocated to a primary instruction from which the secondary instruction depends indicates that there is a hazard related to that primary instruction. Clause 13. The method of any of clauses 1 to 12, wherein a primary instructions is an instructions from which one or more instructions depends and a secondary instructions is an instructions that depends on one or more primary instructions. Clause 14. Computer program code for performing the method of any of clauses 1 to 13. Clause 15. A non-transitory computer readable storage medium having stored thereon computer readable instructions that, when executed at a computer system, cause the computer system to perform the method of any of clauses 1 to 13.

Claims

1. A method of processing instructions in a parallel processing unit comprising a plurality of instruction pipelines, the method comprising: tracking data hazards using a plurality of counters, the plurality of counters comprising a first set of counters associated with high latency data hazards and a second set of counters associated with low latency data hazards; receiving an instruction for execution that indicates (i) whether the instruction is a secondary instruction that depends on at least one primary instruction, and (ii) if the instruction is a secondary instruction, a counter of the plurality of counters associated with each primary instruction from which the instruction depends; determining whether the instruction is a secondary instruction; in response to determining that the instruction is a secondary instruction, analysing the counter associated with each primary instruction from which the instruction depends to determine whether the instruction relates to at least one high latency data hazard that has not been resolved; in response to determining that the instruction relates to at least one high latency data hazard that has not been resolved, causing the instruction to be de-scheduled until each high latency data hazard related to the instruction has been resolved; and in response to determining that the instruction does not relate to at least one high latency data hazard that has not been resolved, forwarding the instruction to a queue preceding an appropriate instruction pipeline of the plurality of instruction pipelines where the instruction stalls until all low latency data hazards related to the instruction have been resolved.
2. The method of claim 1, wherein the instruction further indicates (a) whether the instruction is a primary instruction from which at least one other instruction is dependent, and (b) if the instruction is a primary instruction, a counter of the plurality of counters associated with the primary instruction, and the method further comprising: prior to forwarding the instruction to the queue

preceding the appropriate instruction pipeline, determining whether the instruction is a primary instruction; and in response to determining that the instruction is a primary instruction, causing the counter associated with the primary instruction to be adjusted to indicate that there is a hazard for that primary instruction.

3. The method of claim 2, further comprising, in response to determining that the hazard for a primary instruction has been resolved, causing the counter associated with the primary instruction to be adjusted to indicate that the hazard for the primary instruction has been resolved.

4. The method of claim 2, further comprising, in response to determining that the instruction is a primary instruction, determining whether a wait counter corresponding to the counter associated with the primary instruction indicates that there is at least one secondary instruction waiting on results of that counter; only causing the counter associated with the primary instruction to be adjusted to indicate that there is a hazard for the primary instruction and forwarding the instruction to the queue, in response to determining that the wait counter corresponding to the counter associated with the primary instruction indicates that there are no secondary instructions waiting on the results of that counter; in response to determining that the wait counter corresponding to the counter associated with the primary instruction indicates that there is at least one secondary instruction waiting on the results of the counter, causing the instruction to be de-scheduled until the wait counter corresponding to the counter associated with the primary instruction indicates that there are no secondary instructions waiting on the results of the counter; in response to forwarding a secondary instruction to a queue preceding an appropriate instruction pipeline, causing the value of each wait counter corresponding to the counter associated with a primary instruction from which the secondary instruction depends to be adjusted to indicate that an additional secondary instruction is waiting on the results of the counter; and in response to detecting at a queue that a secondary instruction has been forwarded to an instruction pipeline for execution, cause the value of each wait counter corresponding to the counter associated with a primary instruction from which the secondary instruction depends to be adjusted to indicate that one less secondary instruction is waiting on the results of the counter.

5. The method of claim 1, wherein the instruction further indicates (a) whether the instruction is a primary instruction from which at least one other instruction is dependent, and (b) if the instruction is a primary instruction, a counter of the plurality of counters associated with the primary instruction, and the method further comprising: in response to determining that the instruction is not a secondary instruction, determining whether the instruction is a primary instruction; and in response to determining that the instruction is a primary instruction, causing the counter associated with the primary instruction to be adjusted to indicate that there is a hazard for that primary instruction and forwarding the instruction to the queue preceding the appropriate instruction pipeline.

6. The method of claim 1, wherein the instruction comprises a secondary instruction field that indicates (i) whether the instruction is a secondary instruction and (ii) if the instruction is a secondary instruction, the counter of the plurality of counters associated with each primary instructions from which the instruction depends.

7. The method of claim 6, wherein the secondary instruction field is configured to hold a bit mask, wherein each bit of the bit mask corresponds to a counter of the plurality of counters and when a bit of the bit mask is set, that bit indicates that the instruction is a secondary instruction that is dependent on the primary instruction associated with the corresponding counter.

8. The method of claim 6, wherein the received instruction has been generated by a compiler configured to identify data hazards within a set of related instructions, allocate a counter to each identified data hazard, and generate computer executable instructions that include secondary instruction fields that are configured based on the identifications and counter allocations.

9. The method of claim 1, wherein the instruction comprises a primary instruction field that indicates (a) whether the instruction is a primary instruction from which at least one other

instruction is dependent, and (b) if the instruction is a primary instruction, a counter of the plurality of counters that the primary instruction is associated with.

10. The method of claim 9, wherein the primary instruction field is configured to hold a number and when the number is a predetermined value, the primary instruction field indicates that the instruction is not a primary instruction and when the number is not the predetermined value the primary instruction field indicates that the instruction is a primary instruction and the number identifies the counter associated with the primary instruction.

11. The method of claim 1, wherein causing the instruction to be de-scheduled until each high latency data hazard related to the instruction has been resolved comprises sending a deactivation instruction to a scheduler, the deactivation instruction comprising information identifying the instruction and each high latency data hazard related to the instruction.

12. The method of claim 1, wherein a high latency data hazard is a data hazard associated with a primary instruction that is to be executed by an instruction pipeline of the plurality of instruction pipelines that has a throughput below a predetermined threshold.

13. The method of claim 1, wherein a low latency data hazard is a data hazard associated with a primary instruction that is to be executed by an instruction pipeline of the plurality of instruction pipelines that has a throughput above, or equal to, a predetermined threshold.

14. The method of claim 1, wherein each secondary instruction is to be executed in a different instruction pipeline of the plurality of instruction pipelines from the primary instruction from which it depends.

15. The method of claim 1, wherein the instruction is received at an instruction decoder of the parallel processing unit.

16. A parallel processing unit configured to perform the method as set forth in claim 1.

17. A method of manufacturing, using an integrated circuit manufacturing system, an integrated circuit embodying a parallel processing unit configured to perform the method as set forth in claim 1, comprising inputting into said integrated circuit manufacturing system a computer readable integrated circuit dataset description that describes said parallel processing unit.

18. A non-transitory computer readable storage medium having stored thereon a computer readable dataset description of an integrated circuit that, when processed in an integrated circuit manufacturing system, causes the integrated circuit manufacturing system to manufacture an integrated circuit embodying a parallel processing unit configured to perform the method as set forth in claim 1.

19. An integrated circuit manufacturing system configured to manufacture an integrated circuit embodying a parallel processing unit configured to perform the method as set forth in claim 1.

20. An integrated circuit manufacturing system comprising: a computer readable storage medium having stored thereon a computer readable integrated circuit description that describes a parallel processing unit configured to perform the method as set forth in claim 1; a layout processing system configured to process the integrated circuit description so as to generate a circuit layout description of an integrated circuit embodying the parallel processing unit; and an integrated circuit generation system configured to manufacture the integrated circuit embodying the parallel processing unit according to the circuit layout description.
