

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12405823
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Ashwathnarayan; Sharan et al.

Resource sharing by two or more heterogeneous processing cores

Abstract

Apparatus, systems, and techniques to share memory. In at least one embodiment, a processor comprises one or more circuits to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores.

Inventors: Ashwathnarayan; Sharan (Bangalore, IN), Bhattacharjee; Debalina (Bangalore, IN), Kelur; Ashok (Santa Clara, CA), Parikh; Alok (Bangalore, IN), Kini; Yogesh (Bangalore, IN), Rao; Amit (Bangalore, IN), Paramakuru; Aingarathan (Pickering, CA), Danielson; Kathleen E. (San Francisco, CA), Hettena; Daniel Jonathan (Princeton, NJ), Buzov; Vladislav (San Ramon, CA)

Applicant: NVIDIA Corporation (Santa Clara, CA)

Family ID: 73228460

Assignee: NVIDIA Corporation (Sunnyvale, CA)

Appl. No.: 16/679082

Filed: November 08, 2019

Prior Publication Data

Document Identifier	Publication Date
US 20200364088 A1	Nov. 19, 2020

Foreign Application Priority Data

IN	201911019475	May. 16, 2019
----	--------------	---------------

Publication Classification

Int. Cl.: G06F9/50 (20060101); G06F9/48 (20060101); G06F9/52 (20060101)

U.S. Cl.:

CPC G06F9/5016 (20130101); G06F9/48 (20130101); G06F9/4806 (20130101); G06F9/4843 (20130101); G06F9/4881 (20130101); G06F9/50 (20130101); G06F9/5005 (20130101); G06F9/5027 (20130101); G06F9/5033 (20130101); G06F9/5044 (20130101); G06F9/505 (20130101); G06F9/5088 (20130101); G06F9/52 (20130101); G06F9/522 (20130101); G06F9/526 (20130101);

Field of Classification Search

CPC: G06F (9/48); G06F (9/4806); G06F (9/4843); G06F (9/4881); G06F (9/50); G06F (9/5005); G06F (9/5027); G06F (9/5033); G06F (9/5044); G06F (9/505); G06F (9/5053); G06F (9/5088); G06F (9/52); G06F (9/522); G06F (9/526)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2009/0165014	12/2008	Park	718/105	G06F 9/5088
2009/0222654	12/2008	Hum	713/100	G06F 1/3206
2012/0291040	12/2011	Breternitz	718/104	G06F 9/5083
2014/0101411	12/2013	Sakarda	712/42	G06F 1/329
2014/0240327	12/2013	Lustig	345/505	G06F 9/52
2017/0068571	12/2016	Lu	N/A	G06F 9/5016
2018/0052709	12/2017	Fong	N/A	G06F 9/5044
2018/0074727	12/2017	Zhao	N/A	G06F 3/0656
2018/0143781	12/2017	Greathouse	N/A	G06F 9/451
2019/0312772	12/2018	Zhao	N/A	G06N 3/08
2020/0250585	12/2019	Liu	N/A	G06F 8/447
2020/0334084	12/2019	Jacobson	N/A	G06F 11/3037

OTHER PUBLICATIONS

Damschen et al, "Co-Scheduling on Fused CPU-GPU Architectures With Shared Last Level Caches", IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems, vol. 37, No. 11, Nov. 2018, pp. 2337-2347 (Year: 2018). cited by examiner

Kim et al, "SnuCL: AN OpenCL Framework for Heterogeneous CPU/GPU clusters", ICS'12, Jun. 25-29, 2012, San Servolo Island, Venice, Italy, 2012, ACM, pp. 341-351 (Year: 2012). cited by examiner

Dashti et al, "Analyzing Memory Management Methods on Integrated CPU-GPU Systems", ISMM'17, Jun. 18, 2017, Barcelona, Spain, ACM, pp. 59-69 (Year: 2017). cited by examiner
IEEE Computer Society, "IEEE Standard 754-2008 (Revision of IEEE Standard 754-1985): IEEE Standard for Floating-Point Arithmetic," The Institute of Electrical and Electronics Engineers, Inc., Aug. 29, 2008, 70 pages. cited by applicant

Society of Automotive Engineers On-Road Automated Vehicle Standards Committee "Taxonomy and Definitions for Terms Related to Driving Automation Systems for On-Road Motor Vehicles", Standard No. J3016-201806, dated Jun. 15, 2018. cited by applicant

Society of Automotive Engineers On-Road Automated Vehicle Standards Committee, "Taxonomy

Primary Examiner: Swift; Charles M

Attorney, Agent or Firm: Davis Wright Tremaine LLP

Background/Summary

CROSS REFERENCE TO RELATED APPLICATIONS

(1) This application claims priority to Indian Patent Application No. 201911019475, filed May 16, 2019, entitled “TECHNIQUES FOR STREAMLINED RESOURCE SHARING AND SYNCHRONIZATION ACROSS DISPARATE HARDWARE ENGINES FOR IMPROVED INTEROPERABILITY,” the disclosure of which is incorporated by reference herein in its entirety.

TECHNICAL FIELD

(2) At least one embodiment pertains to facilitating memory sharing between different processor architectures. For example, at least one embodiment, pertains to processors or computing systems used to share memory between different processor architectures according to various novel techniques described herein.

BACKGROUND

(3) Interoperability of memory and other computing resources between different processor architectures and engines can be difficult. Amounts of memory, time, or computing resources used in a system with different processor architectures and engines can be improved.

Description

BRIEF DESCRIPTION OF DRAWINGS

- (1) FIG. 1 illustrates a system that implemented resource allocation and synchronization, according to at least one embodiment;
- (2) FIG. 2 illustrates a diagram that depicts mapping from buffer data object to driver-specific resources, according to at least one embodiment;
- (3) FIG. 3 illustrates a diagram depicting user mode drivers (UMDs) and their allocation semantics, according to at least one embodiment;
- (4) FIG. 4 illustrates a diagram depicting a buffer workflow, according to at least one embodiment;
- (5) FIG. 5 illustrates a diagram depicting buffer attribute validation, according to at least one embodiment;
- (6) FIG. 6 illustrates a diagram depicting importing buffer in parallel computing platform and application programming interface (API) model (e.g., CUDA) external memory interface, according to at least one embodiment;
- (7) FIG. 7 shows an illustrative example of a process to allocate memory, according to at least one embodiment;
- (8) FIG. 8 shows an illustrative example of a process to allocate memory, according to at least one embodiment;
- (9) FIG. 9 illustrates a diagram describing interactions between various objects in an interoperability framework, in accordance with at least one embodiment;
- (10) FIG. 10 illustrates a diagram of semaphore initialization phase, according to at least one embodiment;

(11) FIG. 11 illustrates a diagram of semaphore run phase, according to at least one embodiment;

(12) FIG. 12 illustrates a diagram depicting a graph-based application framework, according to at least one embodiment;

(13) FIG. 13 illustrates a diagram representing an architecture of synchronization, according to at least one embodiment;

(14) FIG. 14 shows an illustrative example of a process to create a synchronization object, according to at least one embodiment;

(15) FIG. 15A illustrates inference and/or training logic, according to at least one embodiment;

(16) FIG. 15B illustrates inference and/or training logic, according to at least one embodiment;

(17) FIG. 16 illustrates training and deployment of a neural network, according to at least one embodiment;

(18) FIG. 17 illustrates an example data center system, according to at least one embodiment;

(19) FIG. 18A illustrates an example of an autonomous vehicle, according to at least one embodiment;

(20) FIG. 18B illustrates an example of camera locations and fields of view for the autonomous vehicle of FIG. 18A, according to at least one embodiment;

(21) FIG. 18C is a block diagram illustrating an example system architecture for the autonomous vehicle of FIG. 18A, according to at least one embodiment;

(22) FIG. 18D is a diagram illustrating a system for communication between cloud-based server(s) and the autonomous vehicle of FIG. 18A, according to at least one embodiment;

(23) FIG. 19 is a block diagram illustrating a computer system, according to at least one embodiment;

(24) FIG. 20 is a block diagram illustrating computer system, according to at least one embodiment;

(25) FIG. 21 illustrates a computer system, according to at least one embodiment;

(26) FIG. 22 illustrates a computer system, according at least one embodiment;

(27) FIG. 23A illustrates a computer system, according to at least one embodiment;

(28) FIG. 23B illustrates a computer system, according to at least one embodiment;

(29) FIG. 23C illustrates a computer system, according to at least one embodiment;

(30) FIG. 23D illustrates a computer system, according to at least one embodiment;

(31) FIGS. 23E and 23F illustrate a shared programming model, according to at least one embodiment;

(32) FIG. 24 illustrates exemplary integrated circuits and associated graphics processors, according to at least one embodiment;

(33) FIGS. 25A-25B illustrate exemplary integrated circuits and associated graphics processors, according to at least one embodiment;

(34) FIGS. 26A-26B illustrate additional exemplary graphics processor logic according to at least one embodiment;

(35) FIG. 27 illustrates a computer system, according to at least one embodiment;

(36) FIG. 28A illustrates a parallel processor, according to at least one embodiment;

(37) FIG. 28B illustrates a partition unit, according to at least one embodiment;

(38) FIG. 28C illustrates a processing cluster, according to at least one embodiment;

(39) FIG. 28D illustrates a graphics multiprocessor, according to at least one embodiment;

(40) FIG. 29 illustrates a multi-graphics processing unit (GPU) system, according to at least one embodiment;

(41) FIG. 30 illustrates a graphics processor, according to at least one embodiment;

(42) FIG. 31 is a block diagram illustrating a processor micro-architecture for a processor, according to at least one embodiment;

(43) FIG. 32 illustrates a deep learning application processor, according to at least one embodiment;

- (44) FIG. **33** is a block diagram illustrating an example neuromorphic processor, according to at least one embodiment;
- (45) FIG. **34** illustrates at least portions of a graphics processor, according to one or more embodiments;
- (46) FIG. **35** illustrates at least portions of a graphics processor, according to one or more embodiments;
- (47) FIG. **36** illustrates at least portions of a graphics processor, according to one or more embodiments;
- (48) FIG. **37** is a block diagram of a graphics processing engine **3710** of a graphics processor in accordance with at least one embodiment.
- (49) FIG. **38** is a block diagram of at least portions of a graphics processor core, according to at least one embodiment;
- (50) FIGS. **39A-39B** illustrate thread execution logic **3900** including an array of processing elements of a graphics processor core according to at least one embodiment
- (51) FIG. **40** illustrates a parallel processing unit (“PPU”), according to at least one embodiment;
- (52) FIG. **41** illustrates a general processing cluster (“GPC”), according to at least one embodiment;
- (53) FIG. **42** illustrates a memory partition unit of a parallel processing unit (“PPU”), according to at least one embodiment;
- (54) FIG. **43** illustrates a streaming multi-processor, according to at least one embodiment;
- (55) FIGS. **44A-44D** illustrate a diagram of unified synchronization, CUDA UMD as writer and reader, according to at least one embodiment; and
- (56) FIGS. **45A-45U** illustrate a diagram of an intra or inter thread use case, according to at least one embodiment.

DETAILED DESCRIPTION

(57) In at least one embodiment, memory allocated for a buffer data object can be imported into a parallel computing platform and application programming interface (API) model (e.g., CUDA). In at least one embodiment, higher level constructs such as image/YUV/tensor can be imported as pointer or arrays according to a parallel computing platform and API model. In at least one embodiment, buffer data object is imported as GPU L2 cache. In at least one embodiment, buffer data object supports both SYSMEM and VIDMEM allocations, wherein SYSMEM may be for access from integrated and discrete GPU engines and VIDMEM is accessible from a discrete GPU (dGPU). In at least one embodiment, buffer data object supports importing memory over process, VM, and chip boundaries. In at least one embodiment, parallel computing platform and application programming interface model allocated memory can be exported as buffer data object. In at least one embodiment, buffer data object allocated memory is interoperable between Tegra platform and x86 platform. In at least one embodiment, buffer data object is a missing link in memory management APIs exposed by different User Mode Drivers (UMDs). In at least one embodiment, data may be shared between multiple engines and consequently multiple UMDs. In at least one embodiment, underlying memory used across engines remains same. In at least one embodiment, there are instances where a particular engine/mode may impose allocation restrictions. In at least one embodiment, one or more allocation restrictions can include one or more of following: for cross-partition mode, allocation should only happen out of Carve-out memory (memory shared between VMs); for display engines, displayable buffers should have specific “Pitch”; parallel computing platform and application programming interface model pointers if used as textures should meet texture alignment requirements; TensorRT should only support Tensor Allocation which is N-dimensional data where N could be between 0 to 8; DLA/PVA engines have constraints like pitch along a particular dimension must be multiple of a certain number “M.”

(58) In at least one embodiment, capabilities of parallel computing platform and application programming interface model external semaphores and parallel computing platform and application

programming interface model streams are enhanced using techniques described herein. In at least one embodiment, a parallel computing platform and application programming interface model stream can wait and signal synchronization object by treating it as a type of external semaphore. In at least one embodiment, parallel computing platform and application programming interface model stream is able to wait for tasks running on a plurality of hardware engines. In at least one embodiment, source sync fence waited upon by a first UMD can be generated by another UMD running on a different engine across a software boundary of thread, process, or VM. In at least one embodiment, a parallel computing platform and application programming interface model stream returns a fence which tracks all tasks currently enqueued in it, which may be similar to EventRecord() except that returned fence represents tasks on stream at that point of time. In at least one embodiment, various types of devices can wait for GPU to finish submitted task (on a stream), completion of which unblocks any device (e.g., across hardware and/or software boundaries) which had enqueued a wait.

(59) In at least one embodiment, synchronization interoperability includes interoperability between parallel computing platform and application programming interface model and synchronization object to allow traditional parallel computing platform and application programming interface model streams to wait for tasks that is outside parallel computing platform and application programming interface model's domain and for other UMDs to natively wait for any tasks enqueued in a parallel computing platform and application programming interface model stream. In at least one embodiment, synchronization interoperability allow applications to gain a finer grain control of efficiently to describe dependencies that spans across hardware and software boundaries. In at least one embodiment, hardware boundary includes one or more of following non-limiting examples: CPU, integrated or discrete GPU, DLA, PVA, ISP, encoder, decode, or an entire chip (e.g., Tegra A, Tegra B on DrivePX2 platforms). In at least one embodiment, software boundaries includes one or more of following non-limiting examples: thread, process, VM.

(60) In at least one embodiment, synchronization object interop is supported as an extension to existing parallel computing platform and application programming interface model interop with external semaphore. In at least one embodiment, parallel computing platform and application programming interface model interops supports externally allocated semaphore, such as Vulkan semaphores. In at least one embodiment, techniques described herein are implemented to support synchronization object interop.

(61) In at least one embodiment, users of synchronization interoperability are abstracted away from internal platform-specific details for portability. In at least one embodiment, sync-primitives are chosen based at last in part on performance to minimize CPU intervention for maintaining dependencies, queries, etc. In at least one embodiment, most performant sync-primitive is chosen. In at least one embodiment, once a synchronization interoperability object is allocated or reserved, it should be reusable across hardware and/or software boundaries. In at least one embodiment, expected functionality is similar to or based on how same event can be used to record task multiple times. In at least one embodiment, reusability of sync interop object reduces resource footprint, thereby reducing memory usage. In at least one embodiment, multiple UMDs are allowed to wait upon a single sync interop object. In at least one embodiment, supporting multi-casting/1:N signaler to waiter relationship is to reduce resource usage and avoid creation of multiple objects that are tracking same task as that of an original interop object, similar to how events can be waited upon in different streams, but different in sense that such waits can happen in different UMDs or processes. In at least one embodiment, synchronization object fits well into higher software abstractions and/or frameworks. In at least one embodiment, synchronization object is to integrate in context of replacement for EGLStream (e.g., NvStream), graph based execution frameworks, user-space scheduler, profiler, and more. In at least one embodiment, synchronization object avoids dynamic memory allocations in critical paths, owing to potential indeterminism of dynamic allocation and possibility of failure (e.g., failure to allocate memory due to system being out of

available memory).

(62) In at least one embodiment, an external entity reserves resources which are shared between signaler & waiter with read/write permissions set appropriately (e.g., reader-writer lock). In at least one embodiment, applications built on one platform with sync interop objects are portable to another platform with little or no change in such applications. In at least one embodiment, portability between x86 platform and Tegra platform with sync interop objects involves few or no changes. In at least one embodiment, synchronization object tracks state of tasks which can be queried. In at least one embodiment, an ability to query for state of task that synchronization object is tracking is supported. In at least one embodiment, task queries help applications and/or schedulers to track completion status, take necessary actions on timeout, and more. In at least one embodiment, CPU/OS backed primitives are supported by synchronization object. In at least one embodiment, device/target are allowed to natively wait upon sync-primitives native to a given OS/CPU (e.g., semaphore in system RAM). In at least one embodiment, for discrete GPUs, semaphores are accessible from video memory (e.g., VIDMEM), which may be opaque to applications (e.g., implementation decision hidden from applications).

(63) FIG. 1 illustrates a system **100** that implemented resource allocation and synchronization, according to at least one embodiment. In at least one embodiment, diagram **100** illustrates resource (memory) **102**; parallel computing platform and application programming interface model array **104**; frame level API library (e.g., NVMedia) Image **106**; parallel computing platform and application programming interface model **108**; and frame level API library (e.g., NVMedia) **110**.

(64) In at least one embodiment, resource **102** is memory shared across two or more UMDs. In at least one embodiment, different UMDs have different attributes and allocation constraints. In at least one embodiment, resource **102** is memory that is allocated based on a set of attributes and list of engines which are collected from each of two or more UMDs. In at least one embodiment, input attributes are processed from all sides to identify a specified engine, constraints are applied to such specified engine, and a buffer is allocated.

(65) In at least one embodiment, sets of attributes and lists of engines are obtained for each of 1 . . . N UMDs. In at least one embodiment, once attribute lists are created for each UMD, lists are to be merged to come up with a set of attributes with which allocation is to be made after applying engine constraints on merged-list. In at least one embodiment, both merging and validation are part of a memory allocation API and not exposed directly to UMDs.

(66) In at least one embodiment, memory allocation is made as per validated attribute list and results in a buffer object handle as well as attribute list handle. In at least one embodiment, a buffer object handle as well as attribute list handle are exposed to applications. In at least one embodiment, buffer object handle is to be mapped to respective UMD VA to be used by application. In at least one embodiment, user can directly invoke query API on attribute list handle to derive final attributes with which allocation was made. In at least one embodiment, UMDs can use object handle to query internal attributes (e.g., RM handle, PageKind) for correct mapping. In at least one embodiment, object handles are passed to UMDs as part of UMD exposed Map API.

(67) In at least one embodiment, resource **102** is memory shared between multiple UMDs wherein resource **102** can be mapped to a parallel computing platform and API model (e.g., CUDA) array **104** in parallel computing platform and application programming interface model **108** UMD and frame level API library (e.g., NVMedia) image **106** in frame level API library (e.g., NVMedia) **110** UMD. In at least one embodiment, access to shared resource **102** can be synchronized between UMDs using a synchronization object that is interoperable between two or more UMDs.

(68) In at least one embodiment, a parallel computing platform and API model refers to an API model that can be used by software developers and software engineers to write code that uses a graphics processing unit (GPU) for general purpose processing. In at least one embodiment, a parallel computing platform and API model is a software layer that gives a programmer or developer direct access to a GPU's virtual instruction set and/or parallel computational elements. In

at least one embodiment, a general-purpose processing unit (GPGPU) refers to an array of GPUs configured to compute highly-parallel operations according to instructions exposed via parallel computing platform and API model. In at least one embodiment, parallel computing platform and API model (e.g., CUDA) is a software platform that can be used for executing compute kernels, which may, based on context, also be referred to as filters. In at least one embodiment, parallel computing platform and API models can be implemented using CUDA, Open Computing Language (OpenCL), DirectCompute, C++ Accelerated Massive Parallelism (C++ AMP), and more. In at least one embodiment described herein, CUDA is used as a non-limiting illustrative example and other parallel computing platform and API models can be used in place of CUDA.

(69) In at least one embodiment, users of synchronization interoperability are abstracted away from internal platform-specific details for portability. In at least one embodiment, sync-primitives are chosen based at least in part on performance to minimize CPU intervention for maintaining dependencies, queries, etc. In at least one embodiment, most performant sync-primitive is chosen. In at least one embodiment, once a synchronization interoperability object is allocated or reserved, it is reusable across hardware and/or software boundaries. In at least one embodiment, reusability of sync interop object reduces resource footprint, thereby reducing memory usage. In at least one embodiment, multiple UMDs are allowed to wait upon a single sync interop object. In at least one embodiment, synchronization object avoids dynamic memory allocations in critical paths, owing to potential indeterminism of dynamic allocation and possibility of failure (e.g., failure to allocate memory due to system being out of available memory).

(70) FIG. 2 illustrates a diagram **200** that depicts mapping from buffer data object to driver-specific resources, according to at least one embodiment. In at least one embodiment, diagram **200** includes allocation semantics **202**, buffer data object **204**, and UMD data objects **206**.

(71) In at least one embodiment, diagram **200** depicts allocation semantics **202**. In at least one embodiment, allocation semantics **202** comprises attributes, compatible types, compatible partitions, and more. In at least one embodiment, attributes are collected from UMDs which are to share memory which is to be allocated in buffer data object **204**. In at least one embodiment, attributes are used at least in part to determine constraints on valid allocations of memory for a buffer data object. In at least one embodiment, compatible types encodes different data types which buffer data object **204** may be interpreted as. In at least one embodiment, a type may be an array, pointer, buffer, texture, tensor, or more. In at least one embodiment, compatible partitions indicates types of partitions which are appropriate for allocation. In at least one embodiment, for cross-partition operation, allocation should be made using carve-out memory (e.g., memory shared between VMs).

(72) In at least one embodiment, allocation semantics **202** are used to determine a manner in which to allocate buffer data object **204**. In at least one embodiment, buffer data object **204** is exposed via a handle which can be interpreted by different UMDs as different UMD-specific data objects. In at least one embodiment, buffer data object **204** encodes additional properties about underlying primitive used for memory allocation including but not limited to memory handles, layouts, properties, sizes, and more. In at least one embodiment, allocation semantics and/or parameters are exposed by an attribute handle which can be queried by UMDs using a buffer API.

(73) In at least one embodiment, a handle to a buffer data object **204** can be interpreted by UMDs to obtain UMD data objects **206**. In at least one embodiment, a SciBuf buffer data object can be interpreted as a parallel computing platform and application programming interface model array or pointer by a first UMD and a frame level API library (e.g., NVMEDIA) image or tensor by a second UMD. In at least one embodiment, UMD data objects are to be used by at least two heterogeneous processing cores. In at least one embodiment, same underlying memory allocation of a buffer data object **204** is shared by at least two heterogeneous processing cores without requiring additional memory copies from one heterogeneous processing core to another.

(74) FIG. 3 illustrates a diagram **300** depicting UMDs and their allocation semantics, according to

at least one embodiment. In at least one embodiment, a buffer supports an allocation of memory to support one or more UMDs illustrated in FIG. 3. In at least one embodiment, diagram **300** includes a set of UMDs **302**. In at least one embodiment, memory allocated to a buffer data object is to be shared across multiple UMDs which may include one or more of parallel computing platform and application programming interface model, frame level API library (e.g., NVMEDIA), and OpenGL. In at least one embodiment, an allocated buffer data object can be mapped to different types of resources **304**. In at least one embodiment, a parallel computing platform and application programming interface model array or parallel computing platform and application programming interface model point, a frame level API library (e.g., NVMEDIA) image or tensor, and an OpenGL texture or buffer are all supported by same underlying allocated buffer data object. In at least one embodiment, a buffer data object can map to a parallel computing platform and application programming interface model array or parallel computing platform and application programming interface model point, a frame level API library (e.g., NVMEDIA) image or tensor, and an cross-language, cross-platform application programming interface for rendering 2D and 3D vector graphics (e.g., OpenGL) texture or buffer. In at least one embodiment, UMDs may have different allocation APIs **306**. In at least one embodiment, memory for a parallel computing platform and application programming interface model array can be allocated using a `ArrayCreate( )` API. In at least one embodiment, memory for a frame level API library (e.g., NVMEDIA) image can be allocated using a frame level API library (e.g., NVMEDIA) `ImageCreate( )` API. In at least one embodiment, memory for an OpenGL texture can be allocated using `glGetTextures( )` API. In at least one embodiment, different UMDs use different memory allocation APIs that have different attributes and/or different constraints on how underlying memory can be allocated. In at least one embodiment, buffer data object memory is allocated in a manner that satisfies constraints of two or more UMDs. In at least one embodiment, attributes and engines from two or more UMDs are collected prior to allocation of buffer data object and merged to determine a manner in which to allocate memory. In at least one embodiment, different UMDs may have different descriptors **308**. In at least one embodiment, SciBuf exposes query API to retrieve parameters with which final allocation was made. In at least one embodiment, descriptors are properties which can be retrieved via API calls. In at least one embodiment, attributes of a buffer data object can be queried. In at least one embodiment, parallel computing platform and application programming interface model array descriptors which can be queried include: width, height, format, number of channels, and any combination thereof.

(75) FIG. 4 illustrates a diagram **400** depicting a buffer workflow, according to at least one embodiment. In at least one embodiment, diagram **400** includes: a first stage **402** for initialization; a second stage **404** for validating and allocating memory; and a third stage **406** for use by UMDs. In at least one embodiment, SciBuf is a type of buffer or memory data object to be shared across two or more UMDs.

(76) In at least one embodiment, diagram **400** includes three stages—a first stage **402** to set attributes and parameters, a second stage to validate and allocate memory, and a third stage to map into UMD space. In at least one embodiment, at a first stage, applications provide inputs. In at least one embodiment, inputs include dimensions of allocations, UMD specific attributes, additional allocation properties (e.g., device, engines, VM), and combinations thereof. In at least one embodiment, corresponding APIs are either exposed directly by SciBuf module or individual UMDs. In at least one embodiment, for parallel computing platform and application programming interface model, applications pass all attributes directly to SciBuf and parallel computing platform and application programming interface model does not expose a public API for same.

(77) In at least one embodiment, a second stage **404** lies entirely with SciBuf, which is responsible for validating if attributes set by different UMDs and generic attributes can result in a valid allocation. In at least one embodiment, if one UMD requests for an allocation on iGPU (Systemem) and another on dGPU (Vidmem), it will result in a validation error. In at least one embodiment, on

success, physical memory is allocated by SciBuf using RM (NvRM) internal interfaces like nvmap APIs and an opaque handle for same is directly exposed to applications. In at least one embodiment, memory domain is selected based on participating UMD specifications in respective attribute lists. In at least one embodiment, SciBuf exposes query API to retrieve parameters with which final allocation was made. In at least one embodiment, parameters which can be retrieved via API calls include: computed pitch; offset; size; and any combination thereof. In at least one embodiment, APIs to retrieve allocation parameters are used to map an allocated buffer into UMD specific objects (VA). Once allocated memory can be imported into VAs in Stage 3. In at least one embodiment, SciBuf exposes a Destroy API to free allocations. In at least one embodiment, for destroying VA mappings, applications can invoke corresponding UMD specific Destroy/Free APIs. (78) In at least one embodiment, a third stage **406** lies with UMDs. In at least one embodiment, a buffer on its own cannot be used directly by any application. In at least one embodiment, buffer can be imported into respective UMD address spaces. In at least one embodiment, for parallel computing platform and application programming interface model, buffer memory can be imported using parallel computing platform and application programming interface model external memory interface. A buffer may refer to a shared buffer or unified buffer.

(79) In at least one embodiment, for teardown, VA mappings as well as actual allocation (SciBuf) may be freed. In at least one embodiment, for freeing UMD VA, UMD specific Destroy APIs can be called by applications. In at least one embodiment, SciBuf exposes API to free backing physical memory. In at least one embodiment, order of free APIs is irrelevant as they result in decrement of refcount and actual memory is freed when refcount becomes zero.

(80) FIG. 5 illustrates a diagram **500** depicting buffer attribute validation, according to at least one embodiment. In at least one embodiment, SciBuf is be a central allocator, whose APIs are exposed to customer applications as well as User Mode Drivers (UMDs). In at least one embodiment, customer applications can specify constraints of all UMDs upfront (e.g., before actual allocation happens). In at least one embodiment, SciBuf ensures allocations can be successful if all constraints can be satisfied, otherwise allocation fails. In at least one embodiment, an allocated buffer can be later on shared with all UMDs whose constraints were specified beforehand. In at least one embodiment, SciBuf supports a single allocation with multiple sharers.

(81) In at least one embodiment, SciBuf receives, as an input, sets of attributes and lists of engines. In at least one embodiment, for each of $1 \dots N$ UMDs, SciBuf receives SciBuf attributes and engine lists. In at least one embodiment, one step in SciBuf allocation is creation of attribute list. In at least one embodiment, AttributeList is an opaque handle externally, internally it is represented as a data structure including a group of attributes with (key, value) pairs. In at least one embodiment, a set of attributes to be set for an SciBuf objects gets decided by its datatype as specified by an application. In at least one embodiment, an application can either choose to use SciBuf public API to create an attribute list or use ones as exposed by UMDs. In at least one embodiment, parallel computing platform and application programming interface model exposes API to set device properties and cache-related information and for other properties like dimensions, application can use SciBuf API directly. In at least one embodiment, UMDs like frame level API library (e.g., NVMedia) expose APIs to set both device and datatype related attributes.

(82) In at least one embodiment, frame level API library (e.g., NVMedia) and parallel computing platform and application programming interface model are APIs at different levels, former being a fixed function driven API recognizes higher level constructs like Image/Tensor, parallel computing platform and application programming interface model, on other hand, being a general purpose driver need not recognize these higher level constructs and accepts datatypes recognized directly by HW or GPU.

(83) In at least one embodiment, once attribute lists are created for each UMD, lists are to be merged to come up with a set of attributes with which allocation is to be made after applying engine constraints on merged-list. In at least one embodiment, both merging and validation are part

of Allocation API exposed and not exposed directly to application. In at least one embodiment, for a cross-process case, allocator process is to invoke Allocate API with attributeLists from all participating processes to come up with an allocation usable across all these processes. In at least one embodiment, medium of communication (IPC) for sharing attribute is left with application. (84) In at least one embodiment, allocation is made as per validated attribute list and results in a buffer object handle as well as attribute list handle. In at least one embodiment, a buffer object handle as well as attribute list handle are exposed to applications. In at least one embodiment, buffer object handle is to be mapped to respective UMD VA to be used by application. In at least one embodiment, user can directly invoke query API on attribute list handle to derive final attributes with which allocation was made. In at least one embodiment, UMDs can use object handle to query internal attributes (RM handle, PageKind, etc.) from SciBuf for correct mapping. In at least one embodiment, object handles are passed to UMDs as part of UMD exposed Map API. In at least one embodiment, for cross-process case, SciBuf allows handle duplication for calling process.

(85) In at least one embodiment, deallocation for a multi-process case happens in multiple steps. In at least one embodiment, UMD references created by current process are to be removed. In at least one embodiment, application are to explicitly invoke SciBuf Free API to unmap object from current process. In at least one embodiment, SciBufFree call is to be invoked by all processes. In at least one embodiment, actual object is to be freed when all process have removed UMD references and local CPU mappings.

(86) FIG. 6 illustrates a diagram 600 depicting importing buffer in parallel computing platform and application programming interface model external memory interface, according to at least one embodiment. In at least one embodiment, Vulkan (or other Graphics API like DX) allocated memory can be imported in parallel computing platform and application programming interface model using APIs provided with parallel computing platform and application programming interface model external memory interface. Once imported, this memory can be mapped in parallel computing platform and application programming interface model specific objects like parallel computing platform and application programming interface model Pointers or parallel computing platform and application programming interface model Arrays. In at least one embodiment, SciBuf allocated memory can be imported into parallel computing platform and application programming interface model using APIs provided with parallel computing platform and application programming interface model external memory interface. In at least one embodiment, import of SciBuf memory uses parallel computing platform and application programming interface model external memory interface with no/minimal changes to existing interfaces. In at least one embodiment, pitch linear and blocklinear memory import is supported. In at least one embodiment, import is allowed across process and VM boundaries. In at least one embodiment, SciBuf allocated in one process/VM can be imported by parallel computing platform and application programming interface model application residing in another process/VM. In at least one embodiment, at time of allocation, SciBuf is aware of its intended usage in parallel computing platform and application programming interface model and device of allocation to impose GPU specific engine restrictions. In at least one embodiment, if applications try to map an already allocated SciBuf into parallel computing platform and application programming interface model domain without setting parallel computing platform and application programming interface model specific attributes, mappings are not guaranteed to succeed. In at least one embodiment, parallel computing platform and application programming interface model Driver imposed constraints (if any) are to be applied before allocation is made. In at least one embodiment, behavior of various parallel computing platform and application programming interface model expose memory allocation APIs (e.g., MemAlloc, MemHostAlloc, ArrayCreate) are achievable with SciBuf_parallel computing platform and application programming interface model interfaces.

(87) In at least one embodiment, SciBuf allows simultaneous access of shared memory from one or

more UMDs and mutually exclusive access to shared buffer is not possible. In at least one embodiment, external semaphore wait/signal APIs can designate hand-off points for parallel computing platform and application programming interface model access to shared buffer and application can invoke external memory APIs to get consistent data.

(88) In at least one embodiment, parallel computing platform and application programming interface model exposes APIs for following functionalities: import an already-allocated SciBuf into parallel computing platform and application programming interface model and delete an imported parallel computing platform and application programming interface model object. In at least one embodiment, externalMemory data type encapsulates memory allocations as a pointer to an opaque struct. In at least one embodiment, name is agnostic and can be extended for non-graphics interops such as SciBuf. In at least one embodiment, externally allocated memory can be imported into parallel computing platform and application programming interface model by providing an appropriate handle. In at least one embodiment, SciBuf handle is identified as `_EXTERNAL_MEMORY_HANDLE_TYPE_SCIBUF`.

(89) In at least one embodiment, a SciBuf handle is associated with an opaque SciBuf object which is returned in a call to `SciBufAllocate( )`. In at least one embodiment, a SciBuf handle holds a reference to underlying SciBuf object which was allocated as per attributes set from one or more drivers—which can include parallel computing platform and application programming interface model. In at least one embodiment, allocation adheres to allocation constraints and dimension requirements of all drivers whose respective `SetSciBufAttrib( )` API was invoked. In at least one embodiment, SciBuf can directly accept parallel computing platform and application programming interface model specific parameter like GPU id. In at least one embodiment, External Memory HandleTypes and Descriptor support SciBuf in following manner:

(90) TABLE-US-00001 Driver Data Structure Change Details external Memory Driver:

```
HandleType typedef enum externalMemoryHandleType_enum {  
    _EXTERNAL_MEMORY_HANDLE_TYPE_OPAQUE_FD = 1, ... .  
    _EXTERNAL_MEMORY_HANDLE_TYPE_SCIBUF = 8, } externalMemoryHandleType;  
_EXTERNAL_MEMORY_HANDLE_DESC Driver: typedef  
struct _EXTERNAL_MEMORY_HANDLE_DESC_st { externalMemoryHandleType type; union {  
    int fd; struct { ... } win32; const void* SciBufObject; } handle; unsigned long long size;  
    unsigned int flags; } _EXTERNAL_MEMORY_HANDLE_DESC;
```

(91) TABLE-US-00002 Runtime Data Structure Change Details

```
ExternalMemoryHandleType_enum Runtime: typedef enum ExternalMemoryHandleType_enum {  
    ExternalMemoryHandleTypeOpaqueFd = 1, ... . ExternalMemoryHandleTypeSciBuf = 8, }  
ExternalMemoryHandleType; ExternalMemoryHandleDesc Runtime: typedef struct  
ExternalMemoryHandleDesc_st { ExternalMemoryHandleType type; union { int fd; struct { ... . }  
    win32; const void* SciBufObject; } handle; unsigned long long size; unsigned int flags; }  
ExternalMemoryHandleDesc;
```

(92) In at least one embodiment, an import API imports an externally-allocated memory object and returns a handle to it. In at least one embodiment, properties of a handle are defined in `_EXTERNAL_MEMORY_HANDLE_DESC`. In at least one embodiment, if handle type is `_EXTERNAL_MEMORY_HANDLE_TYPE_SCIBUF`, `_EXTERNAL_MEMORY_HANDLE_DESC::handle::SciBufhandle::resource` is to be NON NULL and represent a valid SciBuf object. In at least one embodiment, ownership of a SciBuf object is not transferred to parallel computing platform and application programming interface model driver after import operation and remains shared with other drivers who import that SciBuf object in their own address space. In at least one embodiment, proper synchronization and cache operations may be performed by application to avoid overwrites, stale data and undefined behaviors.

(93) In at least one embodiment, a device pointer or a parallel computing platform and application

programming interface model array can be obtained from an external memory allocation by specifying offset and size within previously imported external memory handle. In at least one embodiment, offset and size are to be aligned appropriately and can be queried directly from SciBuf object using NvMem APIs. In at least one embodiment, specifying any other offset and size results in undefined behavior. In at least one embodiment, mapping two buffers whose ranges overlap in external allocation is undefined behavior as they may result in different virtual addresses. In at least one embodiment, for mapping to a parallel computing platform and application programming interface model array, parallel computing platform and application programming interface model Array format is to be specified—mipmapped arrays of level greater than 1 may not be supported with SciBuf handle. In at least one embodiment, once mapped as a parallel computing platform and application programming interface model object, applications can use pointers/arrays as regular ones and perform parallel computing platform and application programming interface model operations like memcpy or memset and pass to parallel computing platform and application programming interface model kernels.

(94) In at least one embodiment, DestroyExternalMemory() API destroys a specified external memory object. In at least one embodiment, existing buffers and parallel computing platform and application programming interface model mipmapped arrays mapped onto a destroyed object are to no longer be used and are to be explicitly freed using MemFree and ArrayDestroy, respectively. In at least one embodiment, once external memory is destroyed, no more mappings should be possible. In at least one embodiment, an application may invoke SciBuf API to free SciBuf object once external memory object is destroyed.

(95) FIG. 7 shows an illustrative example of a process **700** to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with said at least two heterogeneous processing cores, in accordance with at least one embodiment. In at least one embodiment, some or all of process **700** (or any other processes described herein, or variations and/or combinations thereof) is performed under control of one or more computer systems configured with computer-executable instructions and may be implemented as code (e.g., computer-executable instructions, one or more computer programs, or one or more applications) executing collectively on one or more processors, by hardware, software, or combinations thereof. Code, in at least one embodiment, is stored on a computer-readable storage medium in form of a computer program comprising a plurality of computer-readable instructions executable by one or more processors. A computer-readable storage medium, in at least one embodiment, is a non-transitory computer-readable medium. In at least one embodiment, at least some computer-readable instructions usable to perform process **700** are not stored solely using transitory signals (e.g., a propagating transient electric or electromagnetic transmission). A non-transitory computer-readable medium does not necessarily include non-transitory data storage circuitry (e.g., buffers, caches, and queues) within transceivers of transitory signals. In at least one embodiment, process **700** is performed at least in part on a computer system such as those described elsewhere in this disclosure. In at least one embodiment, techniques described in connection with FIG. **8** are utilized in connection with process **700**.

(96) In at least one embodiment, process **700** is implemented by a computer system storing executable instructions that, as a result of execution by one or more processors, obtain **702** one or more attributes associated with at least two heterogeneous processing cores. In at least one embodiment, a heterogeneous processing core described in connection with process **700** is in accordance with those discussed in FIG. **35**. In at least one embodiment, processor cores are homogenous cores executing a common instruction set architecture. In at least one embodiment, processor cores are heterogeneous in terms of instruction set architecture (ISA), where one or more of processor cores execute a common instruction set, while one or more other cores of processor cores executes a subset of a common instruction set or a different instruction set. In at least one

embodiment, processor cores are heterogeneous in terms of microarchitecture, where one or more cores having a relatively higher power consumption couple with one or more power cores having a lower power consumption. In at least one embodiment, processors to implement process **700** can be implemented on one or more chips or as an SoC integrated circuit.

(97) In at least one embodiment, at least two heterogeneous processing cores include two or more UMDs. In at least one embodiment, at least two heterogeneous processing cores comprises a central processing unit (CPU) and a graphics processing unit (GPU). In at least one embodiment, at least two heterogeneous processing cores comprises a central processing unit (CPU) and a graphics processing unit (GPU). In at least one embodiment, at least two heterogeneous processing cores comprises different CPUs supporting different instruction set architectures (e.g., ARM and x86). In at least one embodiment, at least two heterogeneous processing cores comprises an accelerator (e.g., programmable vision accelerator). In at least one embodiment, a system obtains one or more attributes associated with at least two at least two heterogeneous processing cores as lists of attributes, wherein each list of attributes indicates a manner in which a respective UMD plans to utilize memory being allocated. In at least one embodiment, attributes indicate a manner in which a UMD will interpret allocated memory. In at least one embodiment, attributes indicate a type of memory to allocate, such as whether memory is to be allocated using system memory (e.g., DRAM) or video memory (e.g., of a discrete GPU). In at least one embodiment, UMDs specify how memory is to be interpreted according to an enumerated list of attribute types such as in following manner:

(98) TABLE-US-00003 typedef enum attrkeyType{ SciBufAttrKeyType_General, SciBufAttrKeyType_RawBuffer, SciBufAttrKeyType_Image, SciBufAttrKeyType_Tensor, SciBufAttrKeyType_ImagePyramid, SciBufAttrKeyType_Array, SciBufAttrKeyType_Max, }SciBufAttrKeyType;

(99) In at least one embodiment, process **700** is implemented by a computer system storing executable instructions that, as a result of execution by one or more processors, allocate **704** memory according to said one or more attributes. In at least one embodiment, one or more attributes of at least two heterogeneous processing cores determine a set of constraints on how memory to be shared is allocated. In at least one embodiment, contradicting sets of constraints result in an allocation failure. In at least one embodiment a first attribute indicates memory is to be allocated using SYSMEM and a second attribute indicates memory is to be allocated using VIDMEM, thereby resulting in contradictory constraints that result in an error. In at least one embodiment, memory is allocated in a manner that to be interpreted as a first data object by a first heterogeneous processing core and to be interpreted as a second object by a second data object by a second heterogeneous processing core. In at least one embodiment, memory is allocated and returned as a handle to an SciBuf data object which can be interpreted by a first UMD as a parallel computing platform and application programming interface model object (e.g., parallel computing platform and application programming interface model pointer or parallel computing platform and application programming interface model array) and by a second UMD as an OpenGL texture. In at least one embodiment, access to allocated memory by multiple UMDs can be coordinated using techniques described elsewhere in this disclosure, such as those discussed in connection with FIGS. **14** and **15**. In at least one embodiment, a UMD calls a memory allocation API that returns access to shared memory via a handle which can be interpreted by different UMDs as different higher-level data objects.

(100) FIG. **8** shows an illustrative example of a process **800** to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with said at least two heterogeneous processing cores, in accordance with at least one embodiment. In at least one embodiment, some or all of process **800** (or any other processes described herein, or variations and/or combinations thereof) is performed under control

of one or more computer systems configured with computer-executable instructions and may be implemented as code (e.g., computer-executable instructions, one or more computer programs, or one or more applications) executing collectively on one or more processors, by hardware, software, or combinations thereof. Code, in at least one embodiment, is stored on a computer-readable storage medium in form of a computer program comprising a plurality of computer-readable instructions executable by one or more processors. A computer-readable storage medium, in at least one embodiment, is a non-transitory computer-readable medium. In at least one embodiment, at least some computer-readable instructions usable to perform process **800** are not stored solely using transitory signals (e.g., a propagating transient electric or electromagnetic transmission). A non-transitory computer-readable medium does not necessarily include non-transitory data storage circuitry (e.g., buffers, caches, and queues) within transceivers of transitory signals. In at least one embodiment, process **800** is performed at least in part on a computer system such as those described elsewhere in this disclosure. In at least one embodiment, techniques described in connection with FIG. 7 are utilized in connection with process **800**.

(101) In at least one embodiment, process **800** is implemented by a computer system storing executable instructions that, as a result of execution by one or more processors, receive a set of attribute lists describing allocation semantics for a plurality of User Mode Drivers (UMDs). In at least one embodiment, a central allocator receives a set of attributes describing allocation semantics for a plurality of UMDs is received as one or more input parameters of an API. In at least one embodiment, applications running on different UMDs can specify a list of attributes prior to allocation of memory that encodes one or more constraints on memory allocation. In at least one embodiment, lists of attributes are merged and used to determine a set of constraints to be satisfied—if no allocation is able to satisfy all constraints on an allocation, such an allocation may fail. In at least one embodiment, an allocated buffer (e.g., memory) can be shared with all UMDs whose constraints were specified beforehand. In at least one embodiment, memory allocation supports a single allocation with multiple sharers. In at least one embodiment, memory allocated (e.g., in process **800**) is shared memory that can be utilized by multiple UMDs. In at least one embodiment, memory allocated (e.g., in process **800**) is unified memory that can be utilized by multiple UMDs.

(102) In at least one embodiment, a system allocating cross-UMD memory receives, as an input to an API, sets of attributes and lists of engines. In at least one embodiment, a system receives attributes and engine lists for each UMD that is to use shared memory. In at least one embodiment, one step in memory allocation is creation of attribute list. In at least one embodiment, `AttributeList` is an opaque handle externally, internally it is represented as a data structure including a group of attributes with (key, value) pairs. In at least one embodiment, a set of attributes to be set for a buffer gets decided by its datatype as specified by an application. In at least one embodiment, an application can either choose to use a buffer public API to create an attribute list or use ones as exposed by UMDs. In at least one embodiment, parallel computing platform and application programming interface model exposes API to set device properties and cache-related information and for other properties like dimensions, application can use buffer APIs directly. In at least one embodiment, UMDs like frame level API library (e.g., NVMedia) expose APIs to set both device and datatype related attributes.

(103) In at least one embodiment, frame level API library (e.g., NVMedia) and parallel computing platform and application programming interface model are APIs at different levels, former being a fixed function driven API recognizing higher level constructs like Image/Tensor, parallel computing platform and application programming interface model, on other hand, being a general purpose driver which need not recognize higher level constructs and accepts datatypes recognized directly by hardware (e.g., GPU). In at least one embodiment, once attribute lists are created for each UMD, a system performing process **800** includes execution instructions to merge **804** attribute lists to determine allocation constraints. In at least one embodiment, lists are to be merged to determine a set of attributes with which allocation is to be made after applying engine constraints

on merged-list. In at least one embodiment, both merging and validation are part of Allocation API exposed and not exposed directly to application. In at least one embodiment, for a cross-process case, allocator process is to invoke Allocate API with attributeLists from all participating processes to come up with an allocation usable across all these processes. In at least one embodiment, medium of communication (IPC) for sharing attribute is left with application.

(104) In at least one embodiment, a system performing process **800** determines whether **806** it is possible to allocate memory according to allocation constraints determined from merged attribute lists. In at least one embodiment, a system will provide **808** an error message if constraints cannot be satisfied. In at least one embodiment, an allocation constraint that cannot be allocated may include contradictory requirements for: type of memory to use for memory allocation (e.g., SYSMEM vs. VIDMEM); size; memory alignment; and more. An error message, in at least one embodiment, is provided as an error code returned by an API.

(105) In at least one embodiment, if memory can be allocated, a system is to allocate **810** memory as per validated attribute list and results in a buffer object handle as well as attribute list handle. In at least one embodiment, a system is to provide **812** a buffer object handle as well as attribute list handle as output parameters of an API. In at least one embodiment, buffer object handle is to be mapped to respective UMD VA to be used by application. In at least one embodiment, a system (e.g., UMD) is to receive **814** a request to query allocation attributes. In at least one embodiment, user can directly invoke query API on attribute list handle to derive final attributes with which allocation was made. In at least one embodiment, UMDs can use object handle to query internal attributes (e.g., RM handle, PageKind) from SciBuf for correct mapping. In at least one embodiment, system is to map **816** parameters into UMD space and provide response to request according to mappings. In at least one embodiment, object handles are passed to UMDs as part of UMD exposed Map API. In at least one embodiment, for cross-process case, buffer allows handle duplication for calling process.

(106) In at least one embodiment, deallocation for a multi-process case happens in multiple steps. In at least one embodiment, UMD references created by a current process attempting deallocation are to be removed. In at least one embodiment, application are to explicitly invoke a buffer memory deallocation API to unmap object from current process. In at least one embodiment, SciBufFree call is to be invoked by all processes. In at least one embodiment, actual object is to be freed when all process have removed UMD references and local CPU mappings.

(107) FIG. **9** illustrates a diagram **900** describing interactions between various objects in an interoperability framework, in accordance with at least one embodiment. In at least one embodiment, diagram **900** summarizes interactions between various objects including but not limited to: device queue **902**; SciSyncFence **904**; and parallel computing platform and application programming interface model stream **906**.

(108) In at least one embodiment, work is submitted to device queue **902**. In at least one embodiment, device queue **902** is a non-parallel computing platform and application programming interface model queue. In at least one embodiment, an async-signal is queued and SciSyncFence **904** is generated from device queue **902**. In at least one embodiment, device queue **902** is made to wait for generated SciSyncFence **904**. In at least one embodiment, a parallel computing platform and application programming interface model task is a kernel on a parallel computing platform and application programming interface model stream. In at least one embodiment, an already created SciSync is imported as an external semaphore into parallel computing platform and application programming interface model. In at least one embodiment, a signal is issued by parallel computing platform and application programming interface model API WaitExternalSemaphoresAsync. In at least one embodiment, SciSyncFence passed in step 5 illustrated in FIG. **9** was initialized from another UMD and parallel computing platform and application programming interface model will initialize it.

(109) In at least one embodiment, a SciSync object is represented by SciSyncObj which is opaque

to application. In order to import an already created SciSync as an external semaphore into parallel computing platform and application programming interface model, a semaphore descriptors has a reference to a SciSyncObj pointer. In at least one embodiment SciSyncObj pointer is implemented based at least in part on:

(110) TABLE-US-00004 // Driver typedef struct

```
_EXTERNAL_SEMAPHORE_HANDLE_DESC_st { externalSemaphoreHandleType type; union  
{ int fd; struct { void *handle; const void *name; } win32; /** * Valid SciSyncObj. Must be non  
NULL */ const void* SciSyncObj; } handle; unsigned int flags; unsigned int reserved[16]; }
```

_EXTERNAL_SEMAPHORE_HANDLE_DESC; // Runtime typedef struct

```
ExternalSemaphoreHandleDesc_st { ExternalMemoryHandleType type; union { int fd; struct {  
void *handle; const void *name; } win32; const void* SciSyncObj; } handle; unsigned long long  
size; unsigned int flags; unsigned int reserved[16]; } ExternalSemaphoreHandleDesc;
```

(111) In at least one embodiment, an implementation of ImportExternalSemaphore() maps resources into parallel computing platform and application programming interface model's address space and those resources can be accessed at time of signal and wait, it will be subsequently freed at time of DestroyExternalSemaphore(). In at least one embodiment, a resource would be mapping of semaphores associated with SciSync into parallel computing platform and application programming interface model's VA so that acquire and release are done over these address range at time of wait and signal respectively.

(112) In at least one embodiment, parallel computing platform and application programming interface model external semaphores supports importing Vulkan & D3D12 semaphores. In at least one embodiment, to differentiate between already supported types and SciSync,

_EXTERNAL_SEMAPHORE_HANDLE_TYPE_SciSync can be implemented as a new type to externalSemaphoreHandleType. In at least one embodiment, this is set by application before importing SciSync via ImportExternalSemaphore(). In at least one embodiment, a same or similar flag is introduced to runtime version of this structure. In at least one embodiment, externalSemaphoreHandleType is implemented in following manner:

(113) TABLE-US-00005 typedef enum extemalSemaphoreHandleType_enum {

```
_EXTERNAL_SEMAPHORE_HANDLE_TYPE_OPAQUE_FD
```

```
_EXTERNAL_SEMAPHORE_HANDLE_TYPE_OPAQUE_WIN32 = 2,
```

```
_EXTERNAL_SEMAPHORE_HANDLE_TYPE_OPAQUE_WIN32_KMT = 3,
```

```
_EXTERNAL_SEMAPHORE_HANDLE_TYPE_D3D12_FENCE = 4, /** * An opaque handle to  
SciSync */ _EXTERNAL_SEMAPHORE_HANDLE_TYPE_SCISYNC = 5 }
```

extemalSemaphoreHandleType;

(114) In at least on embodiment, ExternalSemaphoreHandleType is a corresponding runtime structure to externalSemaphoreHandleType. In at least one embodiment, a runtime structure is implemented in following manner:

(115) TABLE-US-00006 typedef enum ExternalSemaphoreHandleType_enum {

```
ExtemalSemaphoreHandleTypeOpaqueFd = 1, ExtemalSemaphoreHandleTypeOpaqueWin32 = 2,
```

```
ExtemalSemaphoreHandleTypeOpaqueWin32Kmt = 3,
```

```
ExtemalSemaphoreHandleTypeD3D12Fence = 4, ExtemalSemaphoreHandleTypeSciSync = 5 }
```

ExternalSemaphoreHandleType;

(116) In at least one embodiment, SciSync is designed so that UMD are to wait on an SciSyncFence sent to them by a signaler and return an SciSyncFence for potential waiters to wait. In at least one embodiment, WaitExternalSemaphores or SignalExternalSemaphores APIs accept SciSyncFence in which such APIs are to act. In at least one embodiment, Applications are free to choose their own allocators that backs n SciSyncFence. In at least one embodiment, a UMD simply accept a pointer. In at least one embodiment, a pointer to SciSyncFence is added as described below:

(117) TABLE-US-00007 // Driver typedef struct _EXTERNAL_SEMAPHORE_PARAMS_st {

```

struct { struct { unsigned long long value; } fence; unsigned int reserved[16-
sizeof(void*)]; } params; unsigned int flags; unsigned int reserved[16]; }
_EXTERNAL_SEMAPHORE_PARAMS; // Runtime typedef struct ExternalSemaphoreParams_st
{ struct { struct { unsigned long long value; } fence; unsigned int reserved[16-sizeof(void*)]; void
*SciSyncFence; } params; unsigned int flags; unsigned int reserved[16]; }
ExternalSemaphoreParams;

```

(118) In at least one embodiment, implementation of wait unpacks a fence and issues appropriate acquire methods into a stream. In at least one embodiment, implementation of signal issues appropriate release methods into a stream and fills SciSyncFence structure appropriately.

(119) In at least one embodiment, wait and signal operations occur in pair for Vulkan semaphores. In at least some embodiments, restrictions such as those of Vulkan semaphores are not applicable to SciSync semaphores wherein it is valid for a single SciSyncFence to be waited upon concurrently or otherwise by multiple entities, wherein there is 1:N relationship between number of signals and number of waits. In at least one embodiment, SciSync ensures that a wait is enqueued after a signal is enqueued, which is supported by fact that API which performs a wait is to accept a SciSyncFence that is generated when a signal is enqueued. In at least one embodiment, it is undefined behavior for applications to enqueue wait on invalid SciSyncFences. In at least one embodiment, multiple waits on same SciSyncFence can be enqueued in different threads and process and on different hardware engines, which may be possible because SciSyncFence can be passed-by-value across software boundaries.

(120) In at least one embodiment, WaitExternalSemaphoresAsync() is a supported API. In at least one embodiment, WaitExternalSemaphoresAsync() enqueues a wait operation on a set of externally allocated semaphore objects in a specified stream. In at least one embodiment, operations are executed when all prior operations in a stream are completed. In at least one embodiment, semantics of waiting on a semaphore depend on type of object.

(121) In at least one embodiment, applications invoke WaitExternalSemaphoresAsync() by passing a pointer to SciSyncFence as a parameter via _EXTERNAL_SEMAPHORE_PARAMS. In at least one embodiment, implementation of this API extracts backing synchronization primitive from SciSyncFence, enqueues semaphore acquire or syncPoint acquire method into parallel computing platform and application programming interface model stream. In at least one embodiment, subsequent tasks submitted to this stream have a dependency to whichever tasks was being tracked by SciSyncFence. In at least one embodiment, API follows existing parallel computing platform and application programming interface model stream semantics.

(122) In at least one embodiment, if SciSyncAttrList used to create SciSyncObj had flags in DeviceGetSciSyncAttributes set to _SCISYNC_SIGNAL, API returns _ERROR_NOT_SUPPORTED, since application tried to enqueue a signal while it originally intended to only wait. In at least one embodiment, return value from SignalExternalSemaphoresAsync is made in this regard (similar changes can done to runtime API as well):

(123) TABLE-US-00008 result API WaitExternalSemaphoresAsync(const externalSemaphore *extSemArray, const _EXTERNAL_SEMAPHORE_PARAMS *paramsArray, unsigned int numExtSems, stream stream);

(124) In at least one embodiment, if a semaphore object is any one of following types: _EXTERNAL_SEMAPHORE_HANDLE_TYPE_OPAQUE_FD, _EXTERNAL_SEMAPHORE_HANDLE_TYPE_OPAQUE_WIN32, or _EXTERNAL_SEMAPHORE_HANDLE_TYPE_OPAQUE_WIN32_KMT, then waiting on semaphore waits until semaphore reaches a signaled state. In at least one embodiment, a semaphore reaches a singled state and is then reset to an unsigned state. In at least one embodiment, for every signal operation, there is exactly one corresponding wait operation.

(125) In at least one embodiment, if a semaphore object is type

_EXTERNAL_SEMAPHORE_HANDLE_TYPE_D3D12_FENCE, then waiting on a semaphore waits until value of semaphore is greater than or equal to

_EXTERNAL_SEMAPHORE_PARAMS::params::fence::value.

(126) In at least one embodiment, if a semaphore object is type _EXTERNAL_SEMAPHORE_HANDLE_TYPE_SCISYNC and if SciSyncAttrList used to create SciSyncObj had not set flags in::DeviceGetSciSyncAttributes to _SCISYNC_WAIT, then API returns _ERROR_NOT_SUPPORTED.

(127) In at least one embodiment, WaitExternalSemaphoresAsync() accepts one or more parameters as input parameters. In at least one embodiment a parameter extSemArray refers to external semaphores to be waited on. In at least one embodiment, a parameter paramsArray refers to array of semaphore parameters. In at least one embodiment, a parameter numExtSems refers to a number of semaphores to wait on. In at least one embodiment, a parameter stream refers to external semaphores to a stream to enqueue a wait operation in.

(128) In at least one embodiment, WaitExternalSemaphoresAsync() returns a result value as an output. In at least one embodiment, an output may indicate success or failure states that may include, but are not limited by: not initialized, invalid handle, not supported. In at least one embodiment, result supports one or more of following values: ::_SUCCESS, ::_ERROR_NOT_INITIALIZED, ::_ERROR_INVALID_HANDLE, ::_ERROR_NOT_SUPPORTED

(129) In at least one embodiment, SignalExternalSemaphoresAsync() is a supported API. In at least one embodiment, SignalExternalSemaphoresAsync() enqueues a signal operation on a set of externally allocated semaphore objects in a specified stream. In at least one embodiment, operations will be executed when all prior operations in a stream complete.

(130) In at least one embodiment, applications invokes SignalExternalSemaphoresAsync() by passing a pointer to SciSyncFence as a parameter via _EXTERNAL_SEMAPHORE_PARAMS. In at least one embodiment, API enqueues a signal operation in a parallel computing platform and application programming interface model stream. In at least one embodiment, a signal operation is represented as semaphore release or syncPoint release. In at least one embodiment, operands on which sem_rel or syncPt_rel is issued gets written in a SciSyncFence. In at least one embodiment, SciSyncFence will includes either <SyncPoint-ID, thresholdValue> or <Sema-Offset, thresholdValue>; thresholdvalue being value that a particular syncpoint register or semaphore includes when GPU executes a previously enqueued signal operation.

(131) In at least one embodiment, SciSyncFence tracks one or more same GPU tasks on a parallel computing platform and application programming interface model stream that would have been tracked had a EventRecord() been enqueued into that stream. In at least one embodiment, SciSyncFence returned from this API can be used by other UMDs to wait for completion of parallel computing platform and application programming interface model tasks (e.g., by issuing appropriate acquire methods in their device queue) or natively waiting on SciSyncFence from CPU (e.g., using SciSyncWait()). In at least one embodiment, API follows existing parallel computing platform and application programming interface model stream semantics.

(132) In at least one embodiment, if SciSyncAttrList used to create SciSyncObj had flags in DeviceGetSciSyncAttributes set to _SCISYNC_WAIT, API returns _ERROR_NOT_SUPPORTED, since an application tried to enqueue a wait while it originally intended to only signal. In at least one embodiment, return value from SignalExternalSemaphoresAsync is made in this regard (similar changes can be done to runtime API as well):

(133) TABLE-US-00009 result API SignalExternalSemaphoresAsync(const external Semaphore *extSemArray, const _EXTERNAL_SEMAPHORE_PARAMS *paramsArray, unsigned int numExtSems, stream stream);

(134) In at least one embodiment, semantics of signaling a semaphore depend on type of object. In

at least one embodiment, if semaphore object is any of following types:

`_EXTERNAL_SEMAPHORE_HANDLE_TYPE_OPAQUE_FD`,

`_EXTERNAL_SEMAPHORE_HANDLE_TYPE_OPAQUE_WIN32`, or

`_EXTERNAL_SEMAPHORE_HANDLE_TYPE_OPAQUE_WIN32_KMT`, then signaling semaphore sets it to a signaled state.

(135) In at least one embodiment, if semaphore object is of type

`_EXTERNAL_SEMAPHORE_HANDLE_TYPE_D3D12_FENCE`, then semaphore is set to value specified in `_EXTERNAL_SEMAPHORE_PARAMS::params::fence::value`.

(136) In at least one embodiment, if semaphore object is of type

`_EXTERNAL_SEMAPHORE_HANDLE_TYPE_SCISYNC`, and if `SciSyncAttrList` used to create `SciSyncObj` had not set flags in `::DeviceGetSciSyncAttributes` to `_SCISYNC_SIGNAL`, API returns `::_ERROR_NOT_SUPPORTED`.

(137) In at least one embodiment, `SignalExternalSemaphoresAsync()` accepts one or more

parameters as input parameters. In at least one embodiment a parameter `extSemArray` refers to external semaphores to be signaled. In at least one embodiment, a parameter `paramsArray` refers to array of semaphore parameters. In at least one embodiment, a parameter `numExtSems` refers to a number of semaphores to signal. In at least one embodiment, a parameter `stream` refers to stream to enqueue signal operations in.

(138) In at least one embodiment, `SignalExternalSemaphoresAsync()` returns a result value as an output. In at least one embodiment, an output may indicate success or failure states that may include, but are not limited by: not initialized, invalid handle, not supported. In at least one embodiment, result supports one or more of following values: `::_SUCCESS`,

`::_ERROR_NOT_INITIALIZED`, `::_ERROR_INVALID_HANDLE`,

`::_ERROR_NOT_SUPPORTED`

(139) In at least one embodiment, `WaitExternalSemaphoresAsync()` and

`SignalExternalSemaphoresAsync()` have similar behavior as `WaitExternalSemaphoresAsync()` and `SignalExternalSemaphoresAsync()` and these APIs are not explicitly mentioned again for sake of brevity.

(140) In at least one embodiment, `Cuda-SciSync` is used to ensure data consistency for `Cuda-SciBuf`. In at least one embodiment, `Cuda-SciSync` is used to ensure data consistency for `Cuda-SciBuf` because it is to ensure that APIs are functionally correct and performance optimization would require applications to choose opt-in behavior-elaborating further, when `Cuda-SciBuf` is used by an application and cache-ops are not performed by default, novice application might face bugs which are difficult to debug since user-visible data-consistency between GPU and other engines over shared buffer is dependent on timing, amount of current GPU workload, order of data access etc. Cache-ops enqueued by default, during `Signal & Wait`, ensures that such bugs don't arise in first place. However, unnecessary cache-ops performed has a significant performance impact (full iGPU cache invalidation takes ~4 μ s on Xavier & ~60 μ s on Orin) which is detrimental to auto-cases and unnecessary for use-case where `SciSync` is only needed to express control dependencies.

(141) In at least one embodiment, in order to cater to both novice programmers with limited use-cases and expert programmers with perf-critical use-cases, a flag `_NO_MEMSYNC` can be passed as a parameter to wait & signal API and helps an application let a driver know that `Cuda-SciSync` being operated upon is only capturing control-dependency and there is no need to enqueue cache-ops. In at least some embodiments, a corresponding runtime flag to `_NO_MEMSYNC` is `_NO_MEMSYNC`.

(142) In at least one embodiment, a CPU signals a GPU. In at least one embodiment, ability of this interop to allow CPU signaling GPU deserves a separate section because it provides an approach that is different from what parallel computing platform and application programming interface model supports today. In at least one embodiment, this could potentially be an alternative to `StreamWriteWaitValue`, `StreamAddCallback`.

(143) In at least one embodiment, assume two tasks, C and G where C is a CPU bound task and G is a GPU bound task. Ways in which C G dependency can be built include: Calling thread completes C's execution then submits G to GPU via a stream. Calling thread submits C & G, in that order, in a stream. (e.g., using StreamAddCallback) Using StreamWait or WriteValue APIs. Option 1 is not performant, Option 2 requires applications to follow stream semantics. Option 3 was originally introduced to interop with PCIe devices which owned semaphores. As a side effect it can also be used to build CPU-GPU dependencies using CPU owned semaphores. Even though it is most flexible of 3 options, there is a potential of causing deadlocks. In real world applications on Tegra, task C could: Be running in a separate process. Be written by developers, who are not necessarily parallel computing platform and application programming interface model programmers. Involve accessing non-CUDA SW.

(144) In at least one embodiment, SciSync introduces a way for CPU to signal GPU that is more flexible than option 1 & 2 and while being safer (deadlock free) than option 3. In at least one embodiment, SciSync provides a separate API, SciSyncSignal(), which allows CPU to signal a SciSyncFence. In at least one embodiment, a signaled SciSyncFence could be waited upon in a parallel computing platform and application programming interface model stream. In at least one embodiment, a SciSyncFence created from this category of SciSync could either contain an offset into a semaphore pool or a register_id of syncpoint along with value to wait upon. In at least one embodiment, applications own responsibility of calling SciSyncSignal() at appropriate times, failing which behavior is undefined. In at least one embodiment, if SciSyncSignal() arrives prematurely, a CUDA kernel could start accessing same buffer that is currently being accessed from CPU).

(145) In at least one embodiment, Cuda-SciSync interop follows existing CUDA stream semantics so it is perfectly valid to record an enqueued signal (e.g., via SignalExtSemaAsync()) using a CUDA event. In at least one embodiment, a recorded CUDA event can be used as any regular CUDA event (e.g., build dependencies with other CUDA streams, wait/query an event).

(146) FIG. 10 illustrates a diagram 1000 of semaphore initialization phase, according to at least one embodiment. In at least one embodiment, a signaler (e.g., signaling application on a first hardware engine) calls SyncCreate() with configuration information and a pointer to a handle usSync. In at least one embodiment, SyncCreate() is called by application signaler and fulfilled by Sync module 1006 which sets semaphore as primitive type in usSync, allocates pool for semaphore wherein pool size is equal to number of parallel computing platform and application programming interface model channels multiplied by size of each semaphore. In at least one embodiment, allocation return a handle. In at least one embodiment, MemHandle is saved in Sync in a manner that abstracts from application signalers which underlying synchronization primitive is being used. In at least one embodiment, usSync module returns a status code to a signaling application that indicates whether usSync object was successfully allocated.

(147) In at least one embodiment, signaler 1002 sends usSync object to waiter 1006. In at least one embodiment, sending usSync object to waiter 1006 is optional, such as in case of intra-process waiting. In at least one embodiment, waiter and signaler span different hardware engines. In at least one embodiment, signaler calls a CreateEventFromSync() API where UMD_Signaler 1008 maps usSync object to MemHandle in parallel computing platform and application programming interface model and saves mapped address in a SigEvent. In at least one embodiment, a status code is returned to signaler. In at least one embodiment, waiter calls CreateEventFromSync() API where _UMD_Waiter 1010 maps usSync to MemHandle in parallel computing platform and application programming interface model and saves mapped address in WaitEvent. In at least one embodiment, a status code is returned to waiter. In at least one embodiment, initialization phase ends when both signaler and waiter have created CUDA events for signaling and waiting for events and received status codes, thereby indicating completion of initialization phase. In at least one embodiment, events initialized in FIG. 10 are signaled using techniques described in connection with FIG. 11.

(148) FIG. **11** illustrates a diagram **1100** of semaphore run phase, according to at least one embodiment. In at least one embodiment, FIG. **11** is implemented in context of an initialization phase described in connection with FIG. **10**. In at least one embodiment, signaler **1102** submits work on a parallel computing platform and application programming interface model channel. In at least one embodiment, work comprises a parallel computing platform and application programming interface model kernel accessing a stream. In at least one embodiment, a status code is returned in response to submission of work on a parallel computing platform and application programming interface model channel. In at least one embodiment, signaler calls EventRecord() API and references stream on which work was submitted and SigEvent and provides API call to UMD_Signaler (e.g., same from FIG. **10**). In at least one embodiment, signaler calls ExportSyncFenceFromCudaEvent API which is sent to UMD_Signaler **1104** (e.g., same as from FIG. **10**) which encodes parameters for usSync and SigEvent. In at least one embodiment, UMD_Signaler finds parallel computing platform and application programming interface model channel to track from SigEvent, finds semaphores associated to that channel in pool, adds semaphore release based on channel, address, value information, finds correct offset for that address in pool, and composes offset and value tuple to be returned to signaler. In at least one embodiment, signaler receives a usSyncFence.

(149) In at least one embodiment, usSyncFence is sent to waiter **1106**. In at least one embodiment, sending of usSyncFence is optional, such as in case signaler and waiter are intra-process. In at least one embodiment, waiter calls ImportSyncFenceAsCudaEvent() API with usSync, usSyncFence, and WaitEvent to UMD_Waiter **1108**. In at least one embodiment, UMD_Waiter updates WaitEvent's marker with usSyncFence and returns a status code to waiter. In at least one embodiment, waiter calls StreamWaitEvent() with a second stream and WaitEvent. In at least one embodiment, UMD_Waiter calculates address where sema_acq should be done. In at least one embodiment, address is calculated as WaitEvent reference to semaPool plus offset. In at least one embodiment, UMD_Waiter gets parallel computing platform and application programming interface model channel to add sema_acq from event. In at least one embodiment, UMD_Waiter adds sema_acq with channel, address, and value. In at least one embodiment, a status code is returned to UMD_Waiter. In at least one embodiment, waiter submits work to a second channel. In at least one embodiment, waiter receives a status code in response to work submission.

(150) FIG. **12** illustrates a diagram **1200** depicting a graph-based application framework, according to at least one embodiment. In at least one embodiment, FIG. **12** illustrates a framework to describe an application workflow as a directed acyclic graph (DAG) **1202** with a CPU **1204**, CAM **1206**, iGPU **1208**, and dGPU **1210**. In at least one embodiment, DAG **1202** is similar to parallel computing platform and application programming interface model graphs, but is different in that CUDA graphs lie exclusively in parallel computing platform and application programming interface model domain whereas DAG **1202** allows applications to describe their workflows across different hardware engines, each of which could be governed and/or exposed by a different UMD. In at least one embodiment, a DAG, once described, can be submitted to execution multiple times. In at least one embodiment, DAG **1202** is used to describe a computing environment where images from a camera are to be processed on DLA, iGPU, and dGPU. In at least one embodiment, root node (e.g., CPU node) acts as a trigger to start a task or work.

(151) In at least one embodiment, directed-edges of FIG. **12** represent a control dependencies that exist between node pairs. In at least one embodiment, each edge is backed with a SciSyncFence so that each source node of a directed-edge generates a SciSyncFence that can be waited upon by a corresponding destination node. In at least one embodiment, during graph creation, allocate one Sync per node. In at least one embodiment, each directed edge represents a SyncFence with a waiter and a signaler pair. In at least one embodiment, each run instance of a node updates a SyncFence. In at least one embodiment, dependent nodes can wait on one or more SyncFence. In at least one embodiment, Sync remains implicit to node. In at least one embodiment, graph frameworks choose

right backing primitive.

(152) In at least one embodiment, incoming edges of an edge are processed by extracting backing SyncFence, converting SyncFence to UMD specific type (e.g., CudaEvent) and enqueueing dependency on that engine (e.g., StreamWaitEvent). In at least one embodiment, SyncPoint is used to implement waiting and signaling between CPU and CAM. In at least one embodiment, SyncPoint is used to implement waiting and signaling between CAM and iGPU. In at least one embodiment, semaphore is used to implement waiting and signaling between iGPU and dGPU.

(153) In at least one embodiment, a software application uses a camera (e.g., comprising a first heterogeneous processing core) to take or capture images and CUDA (e.g., running on a second heterogeneous processing core) to read or process those images. In at least one embodiment, a software application creates a first-in-first-out (FIFO) queue (e.g., using a buffer) for fences. In at least one embodiment, a camera checks that a FIFO queue is not full and then takes an image, gets a fence, and adds said image and said fence to said FIFO queue. In at least one embodiment, a CUDA application checks that a FIFO queue is not empty and gets, from said FIFO queue, an image and a fence, adds said fence dependency, and processes said image, for example, by launching a kernel on said image. In at least one embodiment, a fence is a type of synchronization primitive, and techniques described herein may utilize any suitable type of synchronization primitive in place of fences, based on context.

(154) FIG. 13 illustrates a diagram 1300 representing an architecture of synchronization, according to at least one embodiment. In at least one embodiment, synchronization is implemented as SciSync to help different UMDs (e.g., running on same or different device) to signal/wait each other. In at least one embodiment, an OpenGL command queue is able to wait for completion of some CUDA kernel enqueued on a parallel computing platform and application programming interface model stream. In at least one embodiment, a synchronization object refers to a unified synchronization object that is used by two or more UMDs to coordinate execution of code and/or access to data. In at least one embodiment, a synchronization object (e.g., SciSync) is used to coordinate execution of a first set of executable instructions on a first UMD with execution of a second set of executable instructions on a second UMD. In at least one embodiment, a synchronization object is used to coordinate access to memory shared by two or more UMDs. In at least one embodiment, a synchronization object is used to cause a first UMD to wait on a second UMD to provide a signal, after which that first UMD accesses a buffer or memory.

(155) In at least one embodiment, SciSync is used to describe complex dependencies across various engines and different platforms. In at least one embodiment, SciSync helps abstract backing sync-primitives understood by communicating UMDs and operating system specific details from applications. Similar to how parallel computing platform and application programming interface model events abstracts over syncpoints, semaphores (e.g., host or device), QMDs and over an array of operating systems. In at least one embodiment, SciSync, with help of communicating UMDs, reserves all resources to be used for its entire lifetime during initialization so as to avoid need for resource resizing during critical/performant paths. In at least one embodiment, SciSync is designed to stay agnostic to whether or not communicating UMDs are running in separate threads, process, or VMs. In at least one embodiment, depending on properties of communicating UMDs, SciSync helps choose right sync-primitive. In at least one embodiment, DLA only understand syncpoints, while iGPU are capable of understanding syncpoints, semaphores—in this case, SciSync helps CUDA choose syncpoint as right primitive since Syncpoints are a least common entity. In at least one embodiment, if two or more communicating UMDs don't have a least common primitive SciSync doesn't provide any alternate path and such a request for interoperability will fail at creation rather than during a critical/performant path.

(156) In at least one embodiment, CudaEvent and SciSync can be correlated in following manner, with differences highlighted below:

(157) TABLE-US-00010 CudaEvent SciSync A mutable object which tracks tasks submitted An

immutable object which can track tasks to CUDA stream. Each record overwrites submitted to any device queue. previously captured state. Each record overwrites previously captured Since it is immutable, overwriting can never state. happen. CudaEvent:ctxMarker. SciSync: SciSyncFence. Doesn't expose ctxMarker which has all Each Signal will return a SciSyncFence, captured states (QMDs, semaphores) which will always represent state captured. (Syncpoint, semaphore, syncFD) ctxMarker has 1-1 relation with CudaEvents. SciSyncFence has N-1 relation with SciSync. Each record updates state in this unique Each Signal will return a new SciSyncFence ctxMarker. for that captured state. ctxMarker holds reference to context-sensitive SciSyncFence only includes information like VAs. context-insensitive information like SyncPoint register IDs, absolute offsets Supports IPC by exposing events as handles. Support IPCs by exposing SciSync and SciSyncFence as a handle and blob of information respectively. (For inter thread) Applications own Applications have similar responsibilities. responsibility of handling race condition while updating and issuing waits upon an event. (For inter process) Applications should pass Applications can pass around handles that around handles using some mechanism, which represent SciSync and blobs of bytes that could be pipes, mmap, etc. represent SciSyncFence. A preferred mechanism is NvMemBuffer, however applications are free to use other means.

(158) In at least one embodiment, SciSyncFence is implemented as a blob of data. In at least one embodiment, applications and parallel computing platform and application programming interface model driver views SciSyncFence as:

(159) TABLE-US-00011 typedef struct { uint8_t payload[48]; } SciSyncFence;

(160) In at least one embodiment, actual definition as seen by SciSync module is as follows:

(161) TABLE-US-00012 typedef struct SciSyncFenceRec { // Backing synchronization primitive type SciSyncPrimitiveType primitiveType; // A given SciSyncFence will be backed by one of these union { NvU64 syncPointId; // Register ID of SyncPoints NvU64 semaOffset; // Offset into predetermined sema pool . . . ; }; // Threshold value (value for waiters to wait upon) NvU64 value; } SciSyncFence;

(162) In at least one embodiment, applications are allowed to allocate as many SciSyncFences as they wish and pass them for parallel computing platform and application programming interface model specific APIs like wait or signal and managing those is entirely application's responsibility. In at least one embodiment, SciSync module provides getter and setter methods for relevant member. Using these internal interface (non-public) parallel computing platform and application programming interface model drivers can extract appropriate sync-primitives to build dependencies. In at least one embodiment, interfaces are provided by SciSync module to help transfer SciSyncFence across threads, processes, VMs. In at least one embodiment, a given SciSyncFence represents a single point in time and to represent a new point in time, applications can ask for a new SciSyncFence (e.g., by enqueueing a Signal operation).

(163) In at least one embodiment, a SciSyncObj is created from a SciSyncAttrList which, similar to SciSyncFence, has a structure that is opaque to applications. In at least one embodiment, an internal definition is as follows:

(164) TABLE-US-00013 typedef struct { SciSyncUmdType umd; //SciSync_UMD_[parallel computing platform and application programming interface model|GL|. . .] // An array of primitive types that this UMD can support SciSyncPrimitiveType primitiveType[128]; . . . } SciSyncAttrList;

In at least one embodiment, primitiveType array of a SciSyncAttrList holds all base sync-primitives that a given device/engine can interpret/understand. In at least one embodiment, sync-primitives include combinations of syncpoints and semaphores.

(165) In at least one embodiment, all UMDs involved in interop express their capabilities in respective SciSyncAttrList data structures and multiple SciSyncAttrList (e.g., from two or more UMDs) are reconciled to arrive at a combination of capabilities that is common to all involved UMDs. If there is no common capabilities across involved UMDs an error is returned during

reconciliation, else reconciled capabilities are passed to SciSyncObjAlloc() In at least one embodiment, scope of SciSyncObj returned by SciSyncObjAlloc() is limited to calling process. In at least one embodiment, to support cross-process/VM interop, SciSync module provides mechanisms to create a blob of data representing a SciSyncObj which can be transferred across process/VM and/or create a new SciSyncObj from a blob of data representing a SciSyncObj.

(166) In at least one embodiment, different techniques are used to abstract interop between In at least one embodiment, for a case where semaphore is a chosen sync-primitive, a parallel computing platform and application programming interface model API can internally query SciSync module for information regarding a semaphore pool which is returned in a data structure such as in following manner:

(167) TABLE-US-00014

```
typedef struct { // Memory handle for pool. (similar to dma_buf)
    NvMemBuf semaPool; // cache properties NvBool gpuCached; NvBool cpuCached; . . . }
SciSyncSemaphoreInfo;
```

(168) In at least one embodiment, when SciSyncObjAlloc() is called, SciSync module allocates physical memory and stores a reference in semPool, as described above. In at least one embodiment, at time of UMD importing SciSync (e.g., parallel computing platform and application programming interface model), a driver creates a virtual mappings for semaPool. In at least one embodiment, SciSyncSemaphoreInfo is not an opaque data structure to applications. In at least one embodiment, size of pool is set as a part of UMDGetSciSyncAttributes() call. In at least one embodiment, in case of parallel computing platform and application programming interface model size is equal to number of active CUDA channels in current parallel computing platform and application programming interface model context.

(169) In at least one embodiment, every SciSyncFence generated out of such SciSyncObj is backed by a semaphore in pool created earlier. In at least one embodiment, SciSyncFence includes an offset into a pool and a value to wait upon which can be queried by parallel computing platform and application programming interface model using getter methods. In at least one embodiment, information regarding expected cacheability is also sent by SciSync module to parallel computing platform and application programming interface model so that parallel computing platform and application programming interface model can take necessary measures while mapping a semaphore pool. In at least one embodiment, SciSync allocates uncached semaphore. In at least one embodiment, SciSync caches a semaphores on platform with improved **10** coherence support.

(170) In at least one embodiment, syncpoints is chosen to be a backing primitive. In at least one embodiment, a SciSyncFence generated is backed by register_ID and a value in that register to wait upon. In at least one embodiment, for this category of SciSyncs no resource creation is needed at time of an UMD registering SciSync. In at least one embodiment, syncpoint registers are not globally readable (e.g., from different VMs), and necessary syncpoint register needed should be reserved at time of creation; in such scenarios, there would be pool of syncpoint registers that system manages solely to facilitate cross VM/chip communication.

(171) FIG. **14** shows an illustrative example of a process **1400** to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with said at least two heterogeneous processing cores, in accordance with at least one embodiment. In at least one embodiment, some or all of process **1400** (or any other processes described herein, or variations and/or combinations thereof) is performed under control of one or more computer systems configured with computer-executable instructions and may be implemented as code (e.g., computer-executable instructions, one or more computer programs, or one or more applications) executing collectively on one or more processors, by hardware, software, or combinations thereof. Code, in at least one embodiment, is stored on a computer-readable storage medium in form of a computer program comprising a plurality of computer-readable instructions executable by one or more processors. A computer-readable storage medium, in at least

one embodiment, is a non-transitory computer-readable medium. In at least one embodiment, at least some computer-readable instructions usable to perform process **1400** are not stored solely using transitory signals (e.g., a propagating transient electric or electromagnetic transmission). A non-transitory computer-readable medium does not necessarily include non-transitory data storage circuitry (e.g., buffers, caches, and queues) within transceivers of transitory signals. In at least one embodiment, process **1400** is performed at least in part on a computer system such as those described elsewhere in this disclosure.

(172) In at least one embodiment, a system is to obtain **1402** one or more attributes associated with how two or more heterogeneous processing cores support coordinating access to shared memory. In at least one embodiment, shared memory may refer to memory to be accessed (e.g., read and/or write) by two or more UMDs. In at least one embodiment, a heterogeneous processing core may refer to those described elsewhere in this disclosure, such as processor cores described in connection with FIG. 35. In at least one embodiment, a graph-based framework is used to determine a manner in which to coordinate access between heterogeneous processing cores. In at least one embodiment, a first UMD signals and a second UMD wait using a first underlying synchronization object and that second UMD may signal and a third UMD wait using a different synchronization object. In at least one embodiment, two or more UMDs provide attribute lists that include which synchronization objects they support.

(173) In at least one embodiment, a system is to determine **1404**, based on one or more attributes, a manner in which to allocate a synchronization object to coordinate access to memory. In at least one embodiment, a sync object is created from one or more attributes which has a structure that is opaque to applications. In at least one embodiment, one or more attributes includes lists of all base sync-primitives that UMDs are able to interpret/understand. In at least one embodiment, all UMDs involved in interop express their capabilities in respective SciSyncAttrList data structures and multiple SciSyncAttrList (e.g., from two or more UMDs) are reconciled to arrive at a combination of capabilities that is common to all involved UMDs. If there is no common capabilities across involved UMDs an error is returned during reconciliation, else reconciled capabilities are passed to SciSyncObjAlloc(). In at least one embodiment, scope of SciSyncObj returned by SciSyncObjAlloc() is limited to calling process. In at least one embodiment, to support cross-process/VM interop, SciSync module provides mechanisms to create a blob of data representing a SciSyncObj which can be transferred across process/VM and/or create a new SciSyncObj from a blob of data representing a SciSyncObj.

(174) Inference and Training Logic

(175) FIG. 15A illustrates inference and/or training logic **1515** used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided below in conjunction with FIGS. 15A and/or 15B.

(176) In at least one embodiment, inference and/or training logic **1515** may include, without limitation, code and/or data storage **1501** to store forward and/or output weight and/or input/output data, and/or other parameters to configure neurons or layers of a neural network trained and/or used for inferencing in aspects of one or more embodiments. In at least one embodiment, training logic **1515** may include, or be coupled to code and/or data storage **1501** to store graph code or other software to control timing and/or order, in which weight and/or other parameter information is to be loaded to configure, logic, including integer and/or floating point units (collectively, arithmetic logic units (ALUs)). In at least one embodiment, code, such as graph code, loads weight or other parameter information into processor ALUs based on an architecture of a neural network to which the code corresponds. In at least one embodiment code and/or data storage **1501** stores weight parameters and/or input/output data of each layer of a neural network trained or used in conjunction with one or more embodiments during forward propagation of input/output data and/or weight parameters during training and/or inferencing using aspects of one or more embodiments. In at least one embodiment, any portion of code and/or data storage **1501** may be included with other

on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory. (177) In at least one embodiment, any portion of code and/or data storage **1501** may be internal or external to one or more processors or other hardware logic devices or circuits. In at least one embodiment, code and/or code and/or data storage **1501** may be cache memory, dynamic randomly addressable memory (“DRAM”), static randomly addressable memory (“SRAM”), non-volatile memory (e.g., Flash memory), or other storage. In at least one embodiment, choice of whether code and/or code and/or data storage **1501** is internal or external to a processor, for example, or comprised of DRAM, SRAM, Flash or some other storage type may depend on available storage on-chip versus off-chip, latency requirements of training and/or inferencing functions being performed, batch size of data used in inferencing and/or training of a neural network, or some combination of these factors.

(178) In at least one embodiment, inference and/or training logic **1515** may include, without limitation, a code and/or data storage **1505** to store backward and/or output weight and/or input/output data corresponding to neurons or layers of a neural network trained and/or used for inferencing in aspects of one or more embodiments. In at least one embodiment, code and/or data storage **1505** stores weight parameters and/or input/output data of each layer of a neural network trained or used in conjunction with one or more embodiments during backward propagation of input/output data and/or weight parameters during training and/or inferencing using aspects of one or more embodiments. In at least one embodiment, training logic **1515** may include, or be coupled to code and/or data storage **1505** to store graph code or other software to control timing and/or order, in which weight and/or other parameter information is to be loaded to configure, logic, including integer and/or floating point units (collectively, arithmetic logic units (ALUs)). In at least one embodiment, code, such as graph code, loads weight or other parameter information into processor ALUs based on an architecture of a neural network to which the code corresponds. In at least one embodiment, any portion of code and/or data storage **1505** may be included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory. In at least one embodiment, any portion of code and/or data storage **1505** may be internal or external to on one or more processors or other hardware logic devices or circuits. In at least one embodiment, code and/or data storage **1505** may be cache memory, DRAM, SRAM, non-volatile memory (e.g., Flash memory), or other storage. In at least one embodiment, choice of whether code and/or data storage **1505** is internal or external to a processor, for example, or comprised of DRAM, SRAM, Flash or some other storage type may depend on available storage on-chip versus off-chip, latency requirements of training and/or inferencing functions being performed, batch size of data used in inferencing and/or training of a neural network, or some combination of these factors.

(179) In at least one embodiment, code and/or data storage **1501** and code and/or data storage **1505** may be separate storage structures. In at least one embodiment, code and/or data storage **1501** and code and/or data storage **1505** may be same storage structure. In at least one embodiment, code and/or data storage **1501** and code and/or data storage **1505** may be partially same storage structure and partially separate storage structures. In at least one embodiment, any portion of code and/or data storage **1501** and code and/or data storage **1505** may be included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory.

(180) In at least one embodiment, inference and/or training logic **1515** may include, without limitation, one or more arithmetic logic unit(s) (“ALU(s)”) **1510**, including integer and/or floating point units, to perform logical and/or mathematical operations based, at least in part on, or indicated by, training and/or inference code (e.g., graph code), a result of which may produce activations (e.g., output values from layers or neurons within a neural network) stored in an activation storage **1520** that are functions of input/output and/or weight parameter data stored in code and/or data storage **1501** and/or code and/or data storage **1505**. In at least one embodiment, activations stored in activation storage **1520** are generated according to linear algebraic and or

matrix-based mathematics performed by ALU(s) **1510** in response to performing instructions or other code, wherein weight values stored in code and/or data storage **1505** and/or data **1501** are used as operands along with other values, such as bias values, gradient information, momentum values, or other parameters or hyperparameters, any or all of which may be stored in code and/or data storage **1505** or code and/or data storage **1501** or another storage on or off-chip.

(181) In at least one embodiment, ALU(s) **1510** are included within one or more processors or other hardware logic devices or circuits, whereas in another embodiment, ALU(s) **1510** may be external to a processor or other hardware logic device or circuit that uses them (e.g., a co-processor). In at least one embodiment, ALUs **1510** may be included within a processor's execution units or otherwise within a bank of ALUs accessible by a processor's execution units either within same processor or distributed between different processors of different types (e.g., central processing units, graphics processing units, fixed function units, etc.). In at least one embodiment, data storage **1501**, code and/or data storage **1505**, and activation storage **1520** may be on same processor or other hardware logic device or circuit, whereas in another embodiment, they may be in different processors or other hardware logic devices or circuits, or some combination of same and different processors or other hardware logic devices or circuits. In at least one embodiment, any portion of activation storage **1520** may be included with other on-chip or off-chip data storage, including a processor's L1, L2, or L3 cache or system memory. Furthermore, inferencing and/or training code may be stored with other code accessible to a processor or other hardware logic or circuit and fetched and/or processed using a processor's fetch, decode, scheduling, execution, retirement and/or other logical circuits.

(182) In at least one embodiment, activation storage **1520** may be cache memory, DRAM, SRAM, non-volatile memory (e.g., Flash memory), or other storage. In at least one embodiment, activation storage **1520** may be completely or partially within or external to one or more processors or other logical circuits. In at least one embodiment, choice of whether activation storage **1520** is internal or external to a processor, for example, or comprised of DRAM, SRAM, Flash or some other storage type may depend on available storage on-chip versus off-chip, latency requirements of training and/or inferencing functions being performed, batch size of data used in inferencing and/or training of a neural network, or some combination of these factors. In at least one embodiment, inference and/or training logic **1515** illustrated in FIG. **15A** may be used in conjunction with an application-specific integrated circuit ("ASIC"), such as Tensorflow® Processing Unit from Google, an inference processing unit (IPU) from Graphcore™, or a Nervana® (e.g., "Lake Crest") processor from Intel Corp. In at least one embodiment, inference and/or training logic **1515** illustrated in FIG. **15A** may be used in conjunction with central processing unit ("CPU") hardware, graphics processing unit ("GPU") hardware or other hardware, such as field programmable gate arrays ("FPGAs").

(183) FIG. **15B** illustrates inference and/or training logic **1515**, according to at least one embodiment various. In at least one embodiment, inference and/or training logic **1515** may include, without limitation, hardware logic in which computational resources are dedicated or otherwise exclusively used in conjunction with weight values or other information corresponding to one or more layers of neurons within a neural network. In at least one embodiment, inference and/or training logic **1515** illustrated in FIG. **15B** may be used in conjunction with an application-specific integrated circuit (ASIC), such as Tensorflow® Processing Unit from Google, an inference processing unit (IPU) from Graphcore™, or a Nervana® (e.g., "Lake Crest") processor from Intel Corp. In at least one embodiment, inference and/or training logic **1515** illustrated in FIG. **15B** may be used in conjunction with central processing unit (CPU) hardware, graphics processing unit (GPU) hardware or other hardware, such as field programmable gate arrays (FPGAs). In at least one embodiment, inference and/or training logic **1515** includes, without limitation, code and/or data storage **1501** and code and/or data storage **1505**, which may be used to store code (e.g., graph code), weight values and/or other information, including bias values, gradient information,

momentum values, and/or other parameter or hyperparameter information. In at least one embodiment illustrated in FIG. 15B, each of code and/or data storage **1501** and code and/or data storage **1505** is associated with a dedicated computational resource, such as computational hardware **1502** and computational hardware **1506**, respectively. In at least one embodiment, each of computational hardware **1502** and computational hardware **1506** comprises one or more ALUs that perform mathematical functions, such as linear algebraic functions, only on information stored in code and/or data storage **1501** and code and/or data storage **1505**, respectively, result of which is stored in activation storage **1520**.

(184) In at least one embodiment, each of code and/or data storage **1501** and **1505** and corresponding computational hardware **1502** and **1506**, respectively, correspond to different layers of a neural network, such that resulting activation from one “storage/computational pair **1501/1502**” of code and/or data storage **1501** and computational hardware **1502** is provided as an input to next “storage/computational pair **1505/1506**” of code and/or data storage **1505** and computational hardware **1506**, in order to mirror conceptual organization of a neural network. In at least one embodiment, each of storage/computational pairs **1501/1502** and **1505/1506** may correspond to more than one neural network layer. In at least one embodiment, additional storage/computation pairs (not shown) subsequent to or in parallel with storage computation pairs **1501/1502** and **1505/1506** may be included in inference and/or training logic **1515**.

(185) Neural Network Training and Deployment

(186) FIG. 16 illustrates training and deployment of a deep neural network, according to at least one embodiment. In at least one embodiment, untrained neural network **91606** is trained using a training dataset **1602**. In at least one embodiment, training framework **1604** is a PyTorch framework, whereas in other embodiments, training framework **1604** is a Tensorflow, Boost, Caffe, Microsoft Cognitive Toolkit/CNTK, MXNet, Chainer, Keras, Deeplearning4j, or other training framework. In at least one embodiment training framework **1604** trains an untrained neural network **1606** and enables it to be trained using processing resources described herein to generate a trained neural network **1608**. In at least one embodiment, weights may be chosen randomly or by pre-training using a deep belief network. In at least one embodiment, training may be performed in either a supervised, partially supervised, or unsupervised manner.

(187) In at least one embodiment, untrained neural network **1606** is trained using supervised learning, wherein training dataset **1602** includes an input paired with a desired output for an input, or where training dataset **1602** includes input having a known output and an output of neural network **1606** is manually graded. In at least one embodiment, untrained neural network **1606** is trained in a supervised manner processes inputs from training dataset **1602** and compares resulting outputs against a set of expected or desired outputs. In at least one embodiment, errors are then propagated back through untrained neural network **1606**. In at least one embodiment, training framework **1604** adjusts weights that control untrained neural network **1606**. In at least one embodiment, training framework **1604** includes tools to monitor how well untrained neural network **1606** is converging towards a model, such as trained neural network **1608**, suitable to generating correct answers, such as in result **1614**, based on known input data, such as new data **1612**. In at least one embodiment, training framework **1604** trains untrained neural network **1606** repeatedly while adjust weights to refine an output of untrained neural network **1606** using a loss function and adjustment algorithm, such as stochastic gradient descent. In at least one embodiment, training framework **1604** trains untrained neural network **1606** until untrained neural network **1606** achieves a desired accuracy. In at least one embodiment, trained neural network **1608** can then be deployed to implement any number of machine learning operations.

(188) In at least one embodiment, untrained neural network **1606** is trained using unsupervised learning, wherein untrained neural network **1606** attempts to train itself using unlabeled data. In at least one embodiment, unsupervised learning training dataset **1602** will include input data without any associated output data or “ground truth” data. In at least one embodiment, untrained neural

network **1606** can learn groupings within training dataset **1602** and can determine how individual inputs are related to untrained dataset **1602**. In at least one embodiment, unsupervised training can be used to generate a self-organizing map, which is a type of trained neural network **1608** capable of performing operations useful in reducing dimensionality of new data **1612**. In at least one embodiment, unsupervised training can also be used to perform anomaly detection, which allows identification of data points in a new dataset **1612** that deviate from normal patterns of new dataset **1612**.

(189) In at least one embodiment, semi-supervised learning may be used, which is a technique in which in training dataset **1602** includes a mix of labeled and unlabeled data. In at least one embodiment, training framework **1604** may be used to perform incremental learning, such as through transferred learning techniques. In at least one embodiment, incremental learning enables trained neural network **1608** to adapt to new data **1612** without forgetting knowledge instilled within network during initial training.

(190) Data Center

(191) FIG. **17** illustrates an example data center **1700**, in which at least one embodiment may be used. In at least one embodiment, data center **1700** includes a data center infrastructure layer **1710**, a framework layer **1720**, a software layer **1730** and an application layer **1740**.

(192) In at least one embodiment, as shown in FIG. **17**, data center infrastructure layer **1710** may include a resource orchestrator **1712**, grouped computing resources **1714**, and node computing resources (“node C.R.s”) **1716(1)-1716(N)**, where “N” represents any whole, positive integer. In at least one embodiment, node C.R.s **1716(1)-1716(N)** may include, but are not limited to, any number of central processing units (“CPUs”) or other processors (including accelerators, field programmable gate arrays (FPGAs), graphics processors, etc.), memory devices (e.g., dynamic read-only memory), storage devices (e.g., solid state or disk drives), network input/output (“NW I/O”) devices, network switches, virtual machines (“VMs”), power modules, and cooling modules, etc. In at least one embodiment, one or more node C.R.s from among node C.R.s **1716(1)-1716(N)** may be a server having one or more of above-mentioned computing resources.

(193) In at least one embodiment, grouped computing resources **1714** may include separate groupings of node C.R.s housed within one or more racks (not shown), or many racks housed in data centers at various geographical locations (also not shown). separate groupings of node C.R.s within grouped computing resources **1714** may include grouped compute, network, memory or storage resources that may be configured or allocated to support one or more workloads. In at least one embodiment, several node C.R.s including CPUs or processors may grouped within one or more racks to provide compute resources to support one or more workloads. In at least one embodiment, one or more racks may also include any number of power modules, cooling modules, and network switches, in any combination.

(194) In at least one embodiment, resource orchestrator **1712** may configure or otherwise control one or more node C.R.s **1716(1)-1716(N)** and/or grouped computing resources **1714**. In at least one embodiment, resource orchestrator **1712** may include a software design infrastructure (“SDI”) management entity for data center **1700**. In at least one embodiment, resource orchestrator may include hardware, software or some combination thereof.

(195) In at least one embodiment, as shown in FIG. **17**, framework layer **1720** includes a job scheduler **1732**, a configuration manager **1734**, a resource manager **1736** and a distributed file system **1738**. In at least one embodiment, framework layer **1720** may include a framework to support software **1732** of software layer **1730** and/or one or more application(s) **1742** of application layer **1740**. In at least one embodiment, software **1732** or application(s) **1742** may respectively include web-based service software or applications, such as those provided by Amazon Web Services, Google Cloud and Microsoft Azure. In at least one embodiment, framework layer **1720** may be, but is not limited to, a type of free and open-source software web application framework such as Apache Spark™ (hereinafter “Spark”) that may utilize distributed file system **1738** for

large-scale data processing (e.g., “big data”). In at least one embodiment, job scheduler **1732** may include a Spark driver to facilitate scheduling of workloads supported by various layers of data center **1700**. In at least one embodiment, configuration manager **1734** may be capable of configuring different layers such as software layer **1730** and framework layer **1720** including Spark and distributed file system **1738** for supporting large-scale data processing. In at least one embodiment, resource manager **1736** may be capable of managing clustered or grouped computing resources mapped to or allocated for support of distributed file system **1738** and job scheduler **1732**. In at least one embodiment, clustered or grouped computing resources may include grouped computing resource **1714** at data center infrastructure layer **1710**. In at least one embodiment, resource manager **1736** may coordinate with resource orchestrator **1712** to manage these mapped or allocated computing resources.

(196) In at least one embodiment, software **1732** included in software layer **1730** may include software used by at least portions of node C.R.s **1716(1)-1716(N)**, grouped computing resources **1714**, and/or distributed file system **1738** of framework layer **1720**. one or more types of software may include, but are not limited to, Internet web page search software, e-mail virus scan software, database software, and streaming video content software.

(197) In at least one embodiment, application(s) **1742** included in application layer **1740** may include one or more types of applications used by at least portions of node C.R.s **1716(1)-1716(N)**, grouped computing resources **1714**, and/or distributed file system **1738** of framework layer **1720**. one or more types of applications may include, but are not limited to, any number of a genomics application, a cognitive compute, and a machine learning application, including training or inferencing software, machine learning framework software (e.g., PyTorch, TensorFlow, Caffe, etc.) or other machine learning applications used in conjunction with one or more embodiments.

(198) In at least one embodiment, any of configuration manager **1734**, resource manager **1736**, and resource orchestrator **1712** may implement any number and type of self-modifying actions based on any amount and type of data acquired in any technically feasible fashion. In at least one embodiment, self-modifying actions may relieve a data center operator of data center **1700** from making possibly bad configuration decisions and possibly avoiding underutilized and/or poor performing portions of a data center.

(199) In at least one embodiment, data center **1700** may include tools, services, software or other resources to train one or more machine learning models or predict or infer information using one or more machine learning models according to one or more embodiments described herein. For example, in at least one embodiment, a machine learning model may be trained by calculating weight parameters according to a neural network architecture using software and computing resources described above with respect to data center **1700**. In at least one embodiment, trained machine learning models corresponding to one or more neural networks may be used to infer or predict information using resources described above with respect to data center **1700** by using weight parameters calculated through one or more training techniques described herein.

(200) In at least one embodiment, data center may use CPUs, application-specific integrated circuits (ASICs), GPUs, FPGAs, or other hardware to perform training and/or inferencing using above-described resources. Moreover, one or more software and/or hardware resources described above may be configured as a service to allow users to train or performing inferencing of information, such as image recognition, speech recognition, or other artificial intelligence services.

(201) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in system FIG. **17** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(202) In at least one embodiment, data center **1700** runs one or more applications using one or more computing resources which include memory storing computer-readable instructions that, as a result of execution, cause one or more processors to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, data center **1700** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Data center **1700** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-16** and **18-43**.

(203) Autonomous Vehicle

(204) FIG. **18A** illustrates an example of an autonomous vehicle **1800**, according to at least one embodiment. In at least one embodiment, autonomous vehicle **1800** (alternatively referred to herein as “vehicle **1800**”) may be, without limitation, a passenger vehicle, such as a car, a truck, a bus, and/or another type of vehicle that accommodates one or more passengers. In at least one embodiment, vehicle **1800** may be a semi-tractor-trailer truck used for hauling cargo. In at least one embodiment, vehicle **1800** may be an airplane, robotic vehicle, or other kind of vehicle.

(205) Autonomous vehicles may be described in terms of automation levels, defined by National Highway Traffic Safety Administration (“NHTSA”), a division of US Department of Transportation, and Society of Automotive Engineers (“SAE”) “Taxonomy and Definitions for Terms Related to Driving Automation Systems for On-Road Motor Vehicles” (e.g., Standard No. J3016-201806, published on Jun. 15, 2018, Standard No. J3016-201609, published on Sep. 30, 2016, and previous and future versions of this standard). In one or more embodiments, vehicle **1800** may be capable of functionality in accordance with one or more of level 1-level 5 of autonomous driving levels. For example, in at least one embodiment, vehicle **1800** may be capable of conditional automation (Level 3), high automation (Level 4), and/or full automation (Level 5), depending on embodiment.

(206) In at least one embodiment, vehicle **1800** may include, without limitation, components such as a chassis, a vehicle body, wheels (e.g., 2, 4, 6, 8, 18, etc.), tires, axles, and other components of a vehicle. In at least one embodiment, vehicle **1800** may include, without limitation, a propulsion system **1850**, such as an internal combustion engine, hybrid electric power plant, an all-electric engine, and/or another propulsion system type. In at least one embodiment, propulsion system **1850** may be connected to a drive train of vehicle **1800**, which may include, without limitation, a transmission, to enable propulsion of vehicle **1800**. In at least one embodiment, propulsion system **1850** may be controlled in response to receiving signals from a throttle/accelerator(s) **1852**.

(207) In at least one embodiment, a steering system **1854**, which may include, without limitation, a steering wheel, is used to steer a vehicle **1800** (e.g., along a desired path or route) when a propulsion system **1850** is operating (e.g., when vehicle is in motion). In at least one embodiment, a steering system **1854** may receive signals from steering actuator(s) **1856**. steering wheel may be optional for full automation (Level 5) functionality. In at least one embodiment, a brake sensor system **1846** may be used to operate vehicle brakes in response to receiving signals from brake actuator(s) **1848** and/or brake sensors.

(208) In at least one embodiment, controller(s) **1836**, which may include, without limitation, one or more system on chips (“SoCs”) (not shown in FIG. **18A**) and/or graphics processing unit(s) (“GPU(s)”), provide signals (e.g., representative of commands) to one or more components and/or systems of vehicle **1800**. For instance, in at least one embodiment, controller(s) **1836** may send signals to operate vehicle brakes via brake actuators **1848**, to operate steering system **1854** via steering actuator(s) **1856**, to operate propulsion system **1850** via throttle/accelerator(s) **1852**. controller(s) **1836** may include one or more onboard (e.g., integrated) computing devices (e.g.,

supercomputers) that process sensor signals, and output operation commands (e.g., signals representing commands) to enable autonomous driving and/or to assist a human driver in driving vehicle **1800**. In at least one embodiment, controller(s) **1836** may include a first controller **1836** for autonomous driving functions, a second controller **1836** for functional safety functions, a third controller **1836** for artificial intelligence functionality (e.g., computer vision), a fourth controller **1836** for infotainment functionality, a fifth controller **1836** for redundancy in emergency conditions, and/or other controllers. In at least one embodiment, a single controller **1836** may handle two or more of above functionalities, two or more controllers **1836** may handle a single functionality, and/or any combination thereof.

(209) In at least one embodiment, controller(s) **1836** provide signals for controlling one or more components and/or systems of vehicle **1800** in response to sensor data received from one or more sensors (e.g., sensor inputs). In at least one embodiment, sensor data may be received from, for example and without limitation, global navigation satellite systems (“GNSS”) sensor(s) **1858** (e.g., Global Positioning System sensor(s)), RADAR sensor(s) **1860**, ultrasonic sensor(s) **1862**, LIDAR sensor(s) **1864**, inertial measurement unit (“IMU”) sensor(s) **1866** (e.g., accelerometer(s), gyroscope(s), magnetic compass(es), magnetometer(s), etc.), microphone(s) **1896**, stereo camera(s) **1868**, wide-view camera(s) **1870** (e.g., fisheye cameras), infrared camera(s) **1872**, surround camera(s) **1874** (e.g., 360 degree cameras), long-range cameras (not shown in FIG. **18A**), mid-range camera(s) (not shown in FIG. **18A**), speed sensor(s) **1844** (e.g., for measuring speed of vehicle **1800**), vibration sensor(s) **1842**, steering sensor(s) **1840**, brake sensor(s) (e.g., as part of brake sensor system **1846**), and/or other sensor types.

(210) In at least one embodiment, one or more of controller(s) **1836** may receive inputs (e.g., represented by input data) from an instrument cluster **1832** of vehicle **1800** and provide outputs (e.g., represented by output data, display data, etc.) via a human-machine interface (“HMI”) display **1834**, an audible annunciator, a loudspeaker, and/or via other components of vehicle **1800**. In at least one embodiment, outputs may include information such as vehicle velocity, speed, time, map data (e.g., a High Definition map (not shown in FIG. **18A**), location data (e.g., vehicle's **1800** location, such as on a map), direction, location of other vehicles (e.g., an occupancy grid), information about objects and status of objects as perceived by controller(s) **1836**, etc. For example, in at least one embodiment, HMI display **1834** may display information about presence of one or more objects (e.g., a street sign, caution sign, traffic light changing, etc.), and/or information about driving maneuvers vehicle has made, is making, or will make (e.g., changing lanes now, taking exit **34B** in two miles, etc.).

(211) In at least one embodiment, vehicle **1800** further includes a network interface **1824** which may use wireless antenna(s) **1826** and/or modem(s) to communicate over one or more networks. For example, in at least one embodiment, network interface **1824** may be capable of communication over Long-Term Evolution (“LTE”), Wideband Code Division Multiple Access (“WCDMA”), Universal Mobile Telecommunications System (“UMTS”), Global System for Mobile communication (“GSM”), IMT-CDMA Multi-Carrier (“CDMA2000”), etc. In at least one embodiment, wireless antenna(s) **1826** may also enable communication between objects in environment (e.g., vehicles, mobile devices, etc.), using local area network(s), such as Bluetooth, Bluetooth Low Energy (“LE”), Z-Wave, ZigBee, etc., and/or low power wide-area network(s) (“LPWANs”), such as LoRaWAN, SigFox, etc.

(212) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in system FIG. **18A** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(213) In at least one embodiment, vehicle **1800** of FIG. **18A** includes memory storing computer-readable executable instruction that, as a result of execution, causes one or more processors of vehicle **1800** of FIG. **18A** to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. One or more embodiments described elsewhere in this disclosure may be utilized in context of vehicle **1800** of FIG. **18A**, such as techniques described in connection with FIGS. **1-17** and **19-43**.

(214) FIG. **18B** illustrates an example of camera locations and fields of view for autonomous vehicle **1800** of FIG. **18A**, according to at least one embodiment. In at least one embodiment, cameras and respective fields of view are one example embodiment and are not intended to be limiting. For instance, in at least one embodiment, additional and/or alternative cameras may be included and/or cameras may be located at different locations on vehicle **1800**.

(215) In at least one embodiment, camera types for cameras may include, but are not limited to, digital cameras that may be adapted for use with components and/or systems of vehicle **1800**. camera(s) may operate at automotive safety integrity level (“ASIL”) B and/or at another ASIL. In at least one embodiment, camera types may be capable of any image capture rate, such as 60 frames per second (fps), 1220 fps, 240 fps, etc., depending on embodiment. In at least one embodiment, cameras may be capable of using rolling shutters, global shutters, another type of shutter, or a combination thereof. In at least one embodiment, color filter array may include a red clear clear clear (“RCCC”) color filter array, a red clear clear blue (“RCCB”) color filter array, a red blue green clear (“RBGC”) color filter array, a Foveon X3 color filter array, a Bayer sensors (“RGGB”) color filter array, a monochrome sensor color filter array, and/or another type of color filter array. In at least one embodiment, clear pixel cameras, such as cameras with an RCCC, an RCCB, and/or an RBGC color filter array, may be used in an effort to increase light sensitivity.

(216) In at least one embodiment, one or more of camera(s) may be used to perform advanced driver assistance systems (“ADAS”) functions (e.g., as part of a redundant or fail-safe design). For example, in at least one embodiment, a Multi-Function Mono Camera may be installed to provide functions including lane departure warning, traffic sign assist and intelligent headlamp control. In at least one embodiment, one or more of camera(s) (e.g., all of cameras) may record and provide image data (e.g., video) simultaneously.

(217) In at least one embodiment, one or more of cameras may be mounted in a mounting assembly, such as a custom designed (three-dimensional (“3D”) printed) assembly, in order to cut out stray light and reflections from within car (e.g., reflections from dashboard reflected in windshield mirrors) which may interfere with camera's image data capture abilities. With reference to wing-mirror mounting assemblies, in at least one embodiment, wing-mirror assemblies may be custom 3D printed so that camera mounting plate matches shape of wing-mirror. In at least one embodiment, camera(s) may be integrated into wing-mirror. For side-view cameras, camera(s) may also be integrated within four pillars at each corner of cab In at least one embodiment.

(218) In at least one embodiment, cameras with a field of view that include portions of environment in front of vehicle **1800** (e.g., front-facing cameras) may be used for surround view, to help identify forward facing paths and obstacles, as well as aid in, with help of one or more of controllers **1836** and/or control SoCs, providing information critical to generating an occupancy grid and/or determining preferred vehicle paths. In at least one embodiment, front-facing cameras may be used to perform many of same ADAS functions as LIDAR, including, without limitation, emergency braking, pedestrian detection, and collision avoidance. In at least one embodiment, front-facing cameras may also be used for ADAS functions and systems including, without limitation, Lane Departure Warnings (“LDW”), Autonomous Cruise Control (“ACC”), and/or other functions such as traffic sign recognition.

(219) In at least one embodiment, a variety of cameras may be used in a front-facing configuration,

including, for example, a monocular camera platform that includes a CMOS (“complementary metal oxide semiconductor”) color imager. In at least one embodiment, wide-view camera **1870** may be used to perceive objects coming into view from periphery (e.g., pedestrians, crossing traffic or bicycles). Although only one wide-view camera **1870** is illustrated in FIG. **18B**, in other embodiments, there may be any number (including zero) of wide-view camera(s) **1870** on vehicle **1800**. In at least one embodiment, any number of long-range camera(s) **1898** (e.g., a long-view stereo camera pair) may be used for depth-based object detection, especially for objects for which a neural network has not yet been trained. In at least one embodiment, long-range camera(s) **1898** may also be used for object detection and classification, as well as basic object tracking.

(220) In at least one embodiment, any number of stereo camera(s) **1868** may also be included in a front-facing configuration. In at least one embodiment, one or more of stereo camera(s) **1868** may include an integrated control unit comprising a scalable processing unit, which may provide a programmable logic (“FPGA”) and a multi-core micro-processor with an integrated Controller Area Network (“CAN”) or Ethernet interface on a single chip. In at least one embodiment, such a unit may be used to generate a 3D map of environment of vehicle **1800**, including a distance estimate for all points in image. In at least one embodiment, one or more of stereo camera(s) **1868** may include, without limitation, compact stereo vision sensor(s) that may include, without limitation, two camera lenses (one each on left and right) and an image processing chip that may measure distance from vehicle **1800** to target object and use generated information (e.g., metadata) to activate autonomous emergency braking and lane departure warning functions. In at least one embodiment, other types of stereo camera(s) **1868** may be used in addition to, or alternatively from, those described herein.

(221) In at least one embodiment, cameras with a field of view that include portions of environment to side of vehicle **1800** (e.g., side-view cameras) may be used for surround view, providing information used to create and update occupancy grid, as well as to generate side impact collision warnings. For example, in at least one embodiment, surround camera(s) **1874** (e.g., four surround cameras **1874** as illustrated in FIG. **18B**) could be positioned on vehicle **1800**. surround camera(s) **1874** may include, without limitation, any number and combination of wide-view camera(s) **1870**, fisheye camera(s), 360 degree camera(s), and/or like. For instance, in at least one embodiment, four fisheye cameras may be positioned on front, rear, and sides of vehicle **1800**. In at least one embodiment, vehicle **1800** may use three surround camera(s) **1874** (e.g., left, right, and rear), and may leverage one or more other camera(s) (e.g., a forward-facing camera) as a fourth surround-view camera.

(222) In at least one embodiment, cameras with a field of view that include portions of environment to rear of vehicle **1800** (e.g., rear-view cameras) may be used for park assistance, surround view, rear collision warnings, and creating and updating occupancy grid. In at least one embodiment, a wide variety of cameras may be used including, but not limited to, cameras that are also suitable as a front-facing camera(s) (e.g., long-range cameras **1898** and/or mid-range camera(s) **1876**, stereo camera(s) **1868**), infrared camera(s) **1872**, etc.), as described herein.

(223) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in system FIG. **18B** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(224) In at least one embodiment, vehicle **1800** of FIG. **18B** includes memory storing computer-readable executable instruction that, as a result of execution, causes one or more processors of vehicle **1800** of FIG. **18B** to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application

programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. One or more embodiments described elsewhere in this disclosure may be utilized in context of vehicle **1800** of FIG. **18B**, such as techniques described in connection with FIGS. **1-17** and **19-43**.

(225) FIG. **18C** is a block diagram illustrating an example system architecture for autonomous vehicle **1800** of FIG. **18A**, according to at least one embodiment. In at least one embodiment, each of components, features, and systems of vehicle **1800** in FIG. **18C** are illustrated as being connected via a bus **1802**. In at least one embodiment, bus **1802** may include, without limitation, a CAN data interface (alternatively referred to herein as a “CAN bus”). In at least one embodiment, a CAN may be a network inside vehicle **1800** used to aid in control of various features and functionality of vehicle **1800**, such as actuation of brakes, acceleration, braking, steering, windshield wipers, etc. In at least one embodiment, bus **1802** may be configured to have dozens or even hundreds of nodes, each with its own unique identifier (e.g., a CAN ID). In at least one embodiment, bus **1802** may be read to find steering wheel angle, ground speed, engine revolutions per minute (“RPMs”), button positions, and/or other vehicle status indicators. In at least one embodiment, bus **1802** may be a CAN bus that is ASIL B compliant.

(226) In at least one embodiment, in addition to, or alternatively from CAN, FlexRay and/or Ethernet may be used. In at least one embodiment, there may be any number of busses **1802**, which may include, without limitation, zero or more CAN busses, zero or more FlexRay busses, zero or more Ethernet busses, and/or zero or more other types of busses using a different protocol. In at least one embodiment, two or more busses **1802** may be used to perform different functions, and/or may be used for redundancy. For example, a first bus **1802** may be used for collision avoidance functionality and a second bus **1802** may be used for actuation control. In at least one embodiment, each bus **1802** may communicate with any of components of vehicle **1800**, and two or more busses **1802** may communicate with same components. In at least one embodiment, each of any number of system(s) on chip(s) (“SoC(s)”) **1804**, each of controller(s) **1836**, and/or each computer within vehicle may have access to same input data (e.g., inputs from sensors of vehicle **1800**), and may be connected to a common bus, such CAN bus.

(227) In at least one embodiment, vehicle **1800** may include one or more controller(s) **1836**, such as those described herein with respect to FIG. **18A**. controller(s) **1836** may be used for a variety of functions. In at least one embodiment, controller(s) **1836** may be coupled to any of various other components and systems of vehicle **1800**, and may be used for control of vehicle **1800**, artificial intelligence of vehicle **1800**, infotainment for vehicle **1800**, and/or like.

(228) In at least one embodiment, vehicle **1800** may include any number of SoCs **1804**. Each of SoCs **1804** may include, without limitation, central processing units (“CPU(s)”) **1806**, graphics processing units (“GPU(s)”) **1808**, processor(s) **1810**, cache(s) **1812**, accelerator(s) **1814**, data store(s) **1816**, and/or other components and features not illustrated. In at least one embodiment, SoC(s) **1804** may be used to control vehicle **1800** in a variety of platforms and systems. For example, in at least one embodiment, SoC(s) **1804** may be combined in a system (e.g., system of vehicle **1800**) with a High Definition (“HD”) map **1822** which may obtain map refreshes and/or updates via network interface **1824** from one or more servers (not shown in FIG. **18C**).

(229) In at least one embodiment, CPU(s) **1806** may include a CPU cluster or CPU complex (alternatively referred to herein as a “CCPLEX”). In at least one embodiment, CPU(s) **1806** may include multiple cores and/or level two (“L2”) caches. For instance, in at least one embodiment, CPU(s) **1806** may include eight cores in a coherent multi-processor configuration. In at least one embodiment, CPU(s) **1806** may include four dual-core clusters where each cluster has a dedicated L2 cache (e.g., a 2 MB L2 cache). In at least one embodiment, CPU(s) **1806** (e.g., CPLEX) may be configured to support simultaneous cluster operation enabling any combination of clusters of CPU(s) **1806** to be active at any given time.

(230) In at least one embodiment, one or more of CPU(s) **1806** may implement power management

capabilities that include, without limitation, one or more of following features: individual hardware blocks may be clock-gated automatically when idle to save dynamic power; each core clock may be gated when core is not actively executing instructions due to execution of Wait for Interrupt (“WFI”)/Wait for Event (“WFE”) instructions; each core may be independently power-gated; each core cluster may be independently clock-gated when all cores are clock-gated or power-gated; and/or each core cluster may be independently power-gated when all cores are power-gated. In at least one embodiment, CPU(s) **1806** may further implement an enhanced algorithm for managing power states, where allowed power states and expected wakeup times are specified, and hardware/microcode determines best power state to enter for core, cluster, and CCPLX. In at least one embodiment, processing cores may support simplified power state entry sequences in software with work offloaded to microcode.

(231) In at least one embodiment, GPU(s) **1808** may include an integrated GPU (alternatively referred to herein as an “iGPU”). In at least one embodiment, GPU(s) **1808** may be programmable and may be efficient for parallel workloads. In at least one embodiment, GPU(s) **1808**, in at least one embodiment, may use an enhanced tensor instruction set. In on embodiment, GPU(s) **1808** may include one or more streaming microprocessors, where each streaming microprocessor may include a level one (“L1”) cache (e.g., an L1 cache with at least 96 KB storage capacity), and two or more of streaming microprocessors may share an L2 cache (e.g., an L2 cache with a 512 KB storage capacity). In at least one embodiment, GPU(s) **1808** may include at least eight streaming microprocessors. In at least one embodiment, GPU(s) **1808** may use compute application programming interface(s) (API(s)). In at least one embodiment, GPU(s) **1808** may use one or more parallel computing platforms and/or programming models (e.g., NVIDIA's parallel computing platform and application programming interface model).

(232) In at least one embodiment, one or more of GPU(s) **1808** may be power-optimized for best performance in automotive and embedded use cases. For example, in on embodiment, GPU(s) **1808** could be fabricated on a Fin field-effect transistor (“FinFET”). In at least one embodiment, each streaming microprocessor may incorporate a number of mixed-precision processing cores partitioned into multiple blocks. For example, and without limitation, 64 PF32 cores and 32 PF64 cores could be partitioned into four processing blocks. In at least one embodiment, each processing block could be allocated 16 FP32 cores, 8 FP64 cores, 16 INT32 cores, two mixed-precision NVIDIA TENSOR COREs for deep learning matrix arithmetic, a level zero (“L0”) instruction cache, a warp scheduler, a dispatch unit, and/or a 64 KB register file. In at least one embodiment, streaming microprocessors may include independent parallel integer and floating-point data paths to provide for efficient execution of workloads with a mix of computation and addressing calculations. In at least one embodiment, streaming microprocessors may include independent thread scheduling capability to enable finer-grain synchronization and cooperation between parallel threads. In at least one embodiment, streaming microprocessors may include a combined L1 data cache and shared memory unit in order to improve performance while simplifying programming.

(233) In at least one embodiment, one or more of GPU(s) **1808** may include a high bandwidth memory (“HBM”) and/or a 16 GB HBM2 memory subsystem to provide, in some examples, about 900 GB/second peak memory bandwidth. In at least one embodiment, in addition to, or alternatively from, HBM memory, a synchronous graphics random-access memory (“SGRAM”) may be used, such as a graphics double data rate type five synchronous random-access memory (“GDDR5”).

(234) In at least one embodiment, GPU(s) **1808** may include unified memory technology. In at least one embodiment, address translation services (“ATS”) support may be used to allow GPU(s) **1808** to access CPU(s) **1806** page tables directly. In at least one embodiment, embodiment, when GPU(s) **1808** memory management unit (“MMU”) experiences a miss, an address translation request may be transmitted to CPU(s) **1806**. In response, CPU(s) **1806** may look in its page tables for virtual-to-physical mapping for address and transmits translation back to GPU(s) **1808**, in at least one

embodiment. In at least one embodiment, unified memory technology may allow a single unified virtual address space for memory of both CPU(s) **1806** and GPU(s) **1808**, thereby simplifying GPU(s) **1808** programming and porting of applications to GPU(s) **1808**.

(235) In at least one embodiment, GPU(s) **1808** may include any number of access counters that may keep track of frequency of access of GPU(s) **1808** to memory of other processors. In at least one embodiment, access counter(s) may help ensure that memory pages are moved to physical memory of processor that is accessing pages most frequently, thereby improving efficiency for memory ranges shared between processors.

(236) In at least one embodiment, one or more of SoC(s) **1804** may include any number of cache(s) **1812**, including those described herein. For example, in at least one embodiment, cache(s) **1812** could include a level three (“L3”) cache that is available to both CPU(s) **1806** and GPU(s) **1808** (e.g., that is connected both CPU(s) **1806** and GPU(s) **1808**). In at least one embodiment, cache(s) **1812** may include a write-back cache that may keep track of states of lines, such as by using a cache coherence protocol (e.g., MEI, MESI, MSI, etc.). In at least one embodiment, L3 cache may include 4 MB or more, depending on embodiment, although smaller cache sizes may be used.

(237) In at least one embodiment, one or more of SoC(s) **1804** may include one or more accelerator(s) **1814** (e.g., hardware accelerators, software accelerators, or a combination thereof). In at least one embodiment, SoC(s) **1804** may include a hardware acceleration cluster that may include optimized hardware accelerators and/or large on-chip memory. In at least one embodiment, large on-chip memory (e.g., 4 MB of SRAM), may enable hardware acceleration cluster to accelerate neural networks and other calculations. In at least one embodiment, hardware acceleration cluster may be used to complement GPU(s) **1808** and to off-load some of tasks of GPU(s) **1808** (e.g., to free up more cycles of GPU(s) **1808** for performing other tasks). In at least one embodiment, accelerator(s) **1814** could be used for targeted workloads (e.g., perception, convolutional neural networks (“CNNs”), recurrent neural networks (“RNNs”), etc.) that are stable enough to be amenable to acceleration. In at least one embodiment, a CNN may include a region-based or regional convolutional neural networks (“RCNNs”) and Fast RCNNs (e.g., as used for object detection) or other type of CNN.

(238) In at least one embodiment, accelerator(s) **1814** (e.g., hardware acceleration cluster) may include a deep learning accelerator(s) (“DLA”). DLA(s) may include, without limitation, one or more Tensor processing units (“TPUs”) that may be configured to provide an additional ten trillion operations per second for deep learning applications and inferencing. In at least one embodiment, TPUs may be accelerators configured to, and optimized for, performing image processing functions (e.g., for CNNs, RCNNs, etc.). DLA(s) may further be optimized for a specific set of neural network types and floating point operations, as well as inferencing. In at least one embodiment, design of DLA(s) may provide more performance per millimeter than a typical general-purpose GPU, and typically vastly exceeds performance of a CPU. In at least one embodiment, TPU(s) may perform several functions, including a single-instance convolution function, supporting, for example, INT8, INT16, and FP16 data types for both features and weights, as well as post-processor functions. In at least one embodiment, DLA(s) may quickly and efficiently execute neural networks, especially CNNs, on processed or unprocessed data for any of a variety of functions, including, for example and without limitation: a CNN for object identification and detection using data from camera sensors; a CNN for distance estimation using data from camera sensors; a CNN for emergency vehicle detection and identification and detection using data from microphones **1896**; a CNN for facial recognition and vehicle owner identification using data from camera sensors; and/or a CNN for security and/or safety related events.

(239) In at least one embodiment, DLA(s) may perform any function of GPU(s) **1808**, and by using an inference accelerator, for example, a designer may target either DLA(s) or GPU(s) **1808** for any function. For example, in at least one embodiment, designer may focus processing of CNNs and floating point operations on DLA(s) and leave other functions to GPU(s) **1808** and/or other

accelerator(s) **1814**.

(240) In at least one embodiment, accelerator(s) **1814** (e.g., hardware acceleration cluster) may include a programmable vision accelerator(s) (“PVA”), which may alternatively be referred to herein as a computer vision accelerator. In at least one embodiment, PVA(s) may be designed and configured to accelerate computer vision algorithms for advanced driver assistance system (“ADAS”) **1838**, autonomous driving, augmented reality (“AR”) applications, and/or virtual reality (“VR”) applications. PVA(s) may provide a balance between performance and flexibility. For example, in at least one embodiment, each PVA(s) may include, for example and without limitation, any number of reduced instruction set computer (“RISC”) cores, direct memory access (“DMA”), and/or any number of vector processors.

(241) In at least one embodiment, RISC cores may interact with image sensors (e.g., image sensors of any of cameras described herein), image signal processor(s), and/or like. In at least one embodiment, each of RISC cores may include any amount of memory. In at least one embodiment, RISC cores may use any of a number of protocols, depending on embodiment. In at least one embodiment, RISC cores may execute a real-time operating system (“RTOS”). In at least one embodiment, RISC cores may be implemented using one or more integrated circuit devices, application specific integrated circuits (“ASICs”), and/or memory devices. For example, in at least one embodiment, RISC cores could include an instruction cache and/or a tightly coupled RAM.

(242) In at least one embodiment, DMA may enable components of PVA(s) to access system memory independently of CPU(s) **1806**. In at least one embodiment, DMA may support any number of features used to provide optimization to PVA including, but not limited to, supporting multi-dimensional addressing and/or circular addressing. In at least one embodiment, DMA may support up to six or more dimensions of addressing, which may include, without limitation, block width, block height, block depth, horizontal block stepping, vertical block stepping, and/or depth stepping.

(243) In at least one embodiment, vector processors may be programmable processors that may be designed to efficiently and flexibly execute programming for computer vision algorithms and provide signal processing capabilities. In at least one embodiment, PVA may include a PVA core and two vector processing subsystem partitions. In at least one embodiment, PVA core may include a processor subsystem, DMA engine(s) (e.g., two DMA engines), and/or other peripherals. In at least one embodiment, vector processing subsystem may operate as primary processing engine of PVA, and may include a vector processing unit (“VPU”), an instruction cache, and/or vector memory (e.g., “VMEM”). In at least one embodiment, VPU core may include a digital signal processor such as, for example, a single instruction, multiple data (“SIMD”), very long instruction word (“VLIW”) digital signal processor. In at least one embodiment, a combination of SIMD and VLIW may enhance throughput and speed.

(244) In at least one embodiment, each of vector processors may include an instruction cache and may be coupled to dedicated memory. As a result, in at least one embodiment, each of vector processors may be configured to execute independently of other vector processors. In at least one embodiment, vector processors that are included in a particular PVA may be configured to employ data parallelism. For instance, in at least one embodiment, plurality of vector processors included in a single PVA may execute same computer vision algorithm, but on different regions of an image. In at least one embodiment, vector processors included in a particular PVA may simultaneously execute different computer vision algorithms, on same image, or even execute different algorithms on sequential images or portions of an image. In at least one embodiment, among other things, any number of PVAs may be included in hardware acceleration cluster and any number of vector processors may be included in each of PVAs. In at least one embodiment, PVA(s) may include additional error correcting code (“ECC”) memory, to enhance overall system safety.

(245) In at least one embodiment, accelerator(s) **1814** (e.g., hardware acceleration cluster) may include a computer vision network on-chip and static random-access memory (“SRAM”), for

providing a high-bandwidth, low latency SRAM for accelerator(s) **1814**. In at least one embodiment, on-chip memory may include at least 4 MB SRAM, consisting of, for example and without limitation, eight field-configurable memory blocks, that may be accessible by both PVA and DLA. In at least one embodiment, each pair of memory blocks may include an advanced peripheral bus (“APB”) interface, configuration circuitry, a controller, and a multiplexer. In at least one embodiment, any type of memory may be used. In at least one embodiment, PVA and DLA may access memory via a backbone that provides PVA and DLA with high-speed access to memory. In at least one embodiment, backbone may include a computer vision network on-chip that interconnects PVA and DLA to memory (e.g., using APB).

(246) In at least one embodiment, computer vision network on-chip may include an interface that determines, before transmission of any control signal/address/data, that both PVA and DLA provide ready and valid signals. In at least one embodiment, an interface may provide for separate phases and separate channels for transmitting control signals/addresses/data, as well as burst-type communications for continuous data transfer. In at least one embodiment, an interface may comply with International Organization for Standardization (“ISO”) 26262 or International Electrotechnical Commission (“IEC”) 61508 standards, although other standards and protocols may be used.

(247) In at least one embodiment, one or more of SoC(s) **1804** may include a real-time ray-tracing hardware accelerator. In at least one embodiment, real-time ray-tracing hardware accelerator may be used to quickly and efficiently determine positions and extents of objects (e.g., within a world model), to generate real-time visualization simulations, for RADAR signal interpretation, for sound propagation synthesis and/or analysis, for simulation of SONAR systems, for general wave propagation simulation, for comparison to LIDAR data for purposes of localization and/or other functions, and/or for other uses.

(248) In at least one embodiment, accelerator(s) **1814** (e.g., hardware accelerator cluster) have a wide array of uses for autonomous driving. In at least one embodiment, PVA may be a programmable vision accelerator that may be used for key processing stages in ADAS and autonomous vehicles. In at least one embodiment, PVA's capabilities are a good match for algorithmic domains needing predictable processing, at low power and low latency. In other words, PVA performs well on semi-dense or dense regular computation, even on small data sets, which need predictable run-times with low latency and low power. In at least one embodiment, autonomous vehicles, such as vehicle **1800**, PVAs are designed to run classic computer vision algorithms, as they are efficient at object detection and operating on integer math.

(249) For example, according to at least one embodiment of technology, PVA is used to perform computer stereo vision. In at least one embodiment, semi-global matching-based algorithm may be used in some examples, although this is not intended to be limiting. In at least one embodiment, applications for Level 3-5 autonomous driving use motion estimation/stereo matching on-the-fly (e.g., structure from motion, pedestrian recognition, lane detection, etc.). In at least one embodiment, PVA may perform computer stereo vision function on inputs from two monocular cameras.

(250) In at least one embodiment, PVA may be used to perform dense optical flow. For example, in at least one embodiment, PVA could process raw RADAR data (e.g., using a 4D Fast Fourier Transform) to provide processed RADAR data. In at least one embodiment, PVA is used for time of flight depth processing, by processing raw time of flight data to provide processed time of flight data, for example.

(251) In at least one embodiment, DLA may be used to run any type of network to enhance control and driving safety, including for example and without limitation, a neural network that outputs a measure of confidence for each object detection. In at least one embodiment, confidence may be represented or interpreted as a probability, or as providing a relative “weight” of each detection compared to other detections. In at least one embodiment, confidence enables a system to make further decisions regarding which detections should be considered as true positive detections rather

than false positive detections. For example, in at least one embodiment, a system may set a threshold value for confidence and consider only detections exceeding threshold value as true positive detections. In at least one embodiment in which an automatic emergency braking (“AEB”) system is used, false positive detections would cause vehicle to automatically perform emergency braking, which is obviously undesirable. In at least one embodiment, highly confident detections may be considered as triggers for AEB. In at least one embodiment, DLA may run a neural network for regressing confidence value. In at least one embodiment, neural network may take as its input at least some subset of parameters, such as bounding box dimensions, ground plane estimate obtained (e.g. from another subsystem), output from IMU sensor(s) **1866** that correlates with vehicle **1800** orientation, distance, 3D location estimates of object obtained from neural network and/or other sensors (e.g., LIDAR sensor(s) **1864** or RADAR sensor(s) **1860**), among others.

(252) In at least one embodiment, one or more of SoC(s) **1804** may include data store(s) **1816** (e.g., memory). In at least one embodiment, data store(s) **1816** may be on-chip memory of SoC(s) **1804**, which may store neural networks to be executed on GPU(s) **1808** and/or DLA. In at least one embodiment, data store(s) **1816** may be large enough in capacity to store multiple instances of neural networks for redundancy and safety. In at least one embodiment, data store(s) **1812** may comprise L2 or L3 cache(s).

(253) In at least one embodiment, one or more of SoC(s) **1804** may include any number of processor(s) **1810** (e.g., embedded processors). processor(s) **1810** may include a boot and power management processor that may be a dedicated processor and subsystem to handle boot power and management functions and related security enforcement. In at least one embodiment, boot and power management processor may be a part of SoC(s) **1804** boot sequence and may provide runtime power management services. In at least one embodiment, boot power and management processor may provide clock and voltage programming, assistance in system low power state transitions, management of SoC(s) **1804** thermals and temperature sensors, and/or management of SoC(s) **1804** power states. In at least one embodiment, each temperature sensor may be implemented as a ring-oscillator whose output frequency is proportional to temperature, and SoC(s) **1804** may use ring-oscillators to detect temperatures of CPU(s) **1806**, GPU(s) **1808**, and/or accelerator(s) **1814**. In at least one embodiment, if temperatures are determined to exceed a threshold, then boot and power management processor may enter a temperature fault routine and put SoC(s) **1804** into a lower power state and/or put vehicle **1800** into a chauffeur to safe stop mode (e.g., bring vehicle **1800** to a safe stop).

(254) In at least one embodiment, processor(s) **1810** may further include a set of embedded processors that may serve as an audio processing engine. In at least one embodiment, audio processing engine may be an audio subsystem that enables full hardware support for multi-channel audio over multiple interfaces, and a broad and flexible range of audio I/O interfaces. In at least one embodiment, audio processing engine is a dedicated processor core with a digital signal processor with dedicated RAM.

(255) In at least one embodiment, processor(s) **1810** may further include an always on processor engine that may provide necessary hardware features to support low power sensor management and wake use cases. In at least one embodiment, always on processor engine may include, without limitation, a processor core, a tightly coupled RAM, supporting peripherals (e.g., timers and interrupt controllers), various I/O controller peripherals, and routing logic.

(256) In at least one embodiment, processor(s) **1810** may further include a safety cluster engine that includes, without limitation, a dedicated processor subsystem to handle safety management for automotive applications. In at least one embodiment, safety cluster engine may include, without limitation, two or more processor cores, a tightly coupled RAM, support peripherals (e.g., timers, an interrupt controller, etc.), and/or routing logic. In a safety mode, two or more cores may operate, in at least one embodiment, in a lockstep mode and function as a single core with comparison logic to detect any differences between their operations. In at least one embodiment, processor(s) **1810**

may further include a real-time camera engine that may include, without limitation, a dedicated processor subsystem for handling real-time camera management. In at least one embodiment, processor(s) **1810** may further include a high-dynamic range signal processor that may include, without limitation, an image signal processor that is a hardware engine that is part of camera processing pipeline.

(257) In at least one embodiment, processor(s) **1810** may include a video image compositor that may be a processing block (e.g., implemented on a microprocessor) that implements video post-processing functions needed by a video playback application to produce final image for player window. In at least one embodiment, video image compositor may perform lens distortion correction on wide-view camera(s) **1870**, surround camera(s) **1874**, and/or on in-cabin monitoring camera sensor(s). In at least one embodiment, in-cabin monitoring camera sensor(s) are preferably monitored by a neural network running on another instance of SoC **1804**, configured to identify in cabin events and respond accordingly. In at least one embodiment, an in-cabin system may perform, without limitation, lip reading to activate cellular service and place a phone call, dictate emails, change vehicle's destination, activate or change vehicle's infotainment system and settings, or provide voice-activated web surfing. In at least one embodiment, certain functions are available to driver when vehicle is operating in an autonomous mode and are disabled otherwise.

(258) In at least one embodiment, video image compositor may include enhanced temporal noise reduction for both spatial and temporal noise reduction. For example, in at least one embodiment, where motion occurs in a video, noise reduction weights spatial information appropriately, decreasing weight of information provided by adjacent frames. In at least one embodiment, where an image or portion of an image does not include motion, temporal noise reduction performed by video image compositor may use information from previous image to reduce noise in current image.

(259) In at least one embodiment, video image compositor may also be configured to perform stereo rectification on input stereo lens frames. In at least one embodiment, video image compositor may further be used for user interface composition when operating system desktop is in use, and GPU(s) **1808** are not required to continuously render new surfaces. In at least one embodiment, when GPU(s) **1808** are powered on and active doing 3D rendering, video image compositor may be used to offload GPU(s) **1808** to improve performance and responsiveness.

(260) In at least one embodiment, one or more of SoC(s) **1804** may further include a mobile industry processor interface (“MIPI”) camera serial interface for receiving video and input from cameras, a high-speed interface, and/or a video input block that may be used for camera and related pixel input functions. In at least one embodiment, one or more of SoC(s) **1804** may further include an input/output controller(s) that may be controlled by software and may be used for receiving I/O signals that are uncommitted to a specific role.

(261) In at least one embodiment, one or more of SoC(s) **1804** may further include a broad range of peripheral interfaces to enable communication with peripherals, audio encoders/decoders (“codecs”), power management, and/or other devices. SoC(s) **1804** may be used to process data from cameras (e.g., connected over Gigabit Multimedia Serial Link and Ethernet), sensors (e.g., LIDAR sensor(s) **1864**, RADAR sensor(s) **1860**, etc. that may be connected over Ethernet), data from bus **1802** (e.g., speed of vehicle **1800**, steering wheel position, etc.), data from GNSS sensor(s) **1858** (e.g., connected over Ethernet or CAN bus), etc. In at least one embodiment, one or more of SoC(s) **1804** may further include dedicated high-performance mass storage controllers that may include their own DMA engines, and that may be used to free CPU(s) **1806** from routine data management tasks.

(262) In at least one embodiment, SoC(s) **1804** may be an end-to-end platform with a flexible architecture that spans automation levels 3-5, thereby providing a comprehensive functional safety architecture that leverages and makes efficient use of computer vision and ADAS techniques for diversity and redundancy, provides a platform for a flexible, reliable driving software stack, along

with deep learning tools. In at least one embodiment, SoC(s) **1804** may be faster, more reliable, and even more energy-efficient and space-efficient than conventional systems. For example, in at least one embodiment, accelerator(s) **1814**, when combined with CPU(s) **1806**, GPU(s) **1808**, and data store(s) **1816**, may provide for a fast, efficient platform for level 3-5 autonomous vehicles.

(263) In at least one embodiment, computer vision algorithms may be executed on CPUs, which may be configured using high-level programming language, such as C programming language, to execute a wide variety of processing algorithms across a wide variety of visual data. However, in at least one embodiment, CPUs are oftentimes unable to meet performance requirements of many computer vision applications, such as those related to execution time and power consumption, for example. In at least one embodiment, many CPUs are unable to execute complex object detection algorithms in real-time, which is used in in-vehicle ADAS applications and in practical Level 3-5 autonomous vehicles.

(264) Embodiments described herein allow for multiple neural networks to be performed simultaneously and/or sequentially, and for results to be combined together to enable Level 3-5 autonomous driving functionality. For example, in at least one embodiment, a CNN executing on DLA or discrete GPU (e.g., GPU(s) **1820**) may include text and word recognition, allowing supercomputer to read and understand traffic signs, including signs for which neural network has not been specifically trained. In at least one embodiment, DLA may further include a neural network that is able to identify, interpret, and provide semantic understanding of sign, and to pass that semantic understanding to path planning modules running on CPU Complex.

(265) In at least one embodiment, multiple neural networks may be run simultaneously, as for Level 3, 4, or 5 driving. For example, in at least one embodiment, a warning sign consisting of “Caution: flashing lights indicate icy conditions,” along with an electric light, may be independently or collectively interpreted by several neural networks. In at least one embodiment, sign itself may be identified as a traffic sign by a first deployed neural network (e.g., a neural network that has been trained), text “flashing lights indicate icy conditions” may be interpreted by a second deployed neural network, which informs vehicle's path planning software (preferably executing on CPU Complex) that when flashing lights are detected, icy conditions exist. In at least one embodiment, flashing light may be identified by operating a third deployed neural network over multiple frames, informing vehicle's path-planning software of presence (or absence) of flashing lights. In at least one embodiment, all three neural networks may run simultaneously, such as within DLA and/or on GPU(s) **1808**.

(266) In at least one embodiment, a CNN for facial recognition and vehicle owner identification may use data from camera sensors to identify presence of an authorized driver and/or owner of vehicle **1800**. In at least one embodiment, an always on sensor processing engine may be used to unlock vehicle when owner approaches driver door and turn on lights, and, in security mode, to disable vehicle when owner leaves vehicle. In this way, SoC(s) **1804** provide for security against theft and/or carjacking.

(267) In at least one embodiment, a CNN for emergency vehicle detection and identification may use data from microphones **1896** to detect and identify emergency vehicle sirens. In at least one embodiment, SoC(s) **1804** use CNN for classifying environmental and urban sounds, as well as classifying visual data. In at least one embodiment, CNN running on DLA is trained to identify relative closing speed of emergency vehicle (e.g., by using Doppler effect). In at least one embodiment, CNN may also be trained to identify emergency vehicles specific to local area in which vehicle is operating, as identified by GNSS sensor(s) **1858**. In at least one embodiment, when operating in Europe, CNN will seek to detect European sirens, and when in United States CNN will seek to identify only North American sirens. In at least one embodiment, once an emergency vehicle is detected, a control program may be used to execute an emergency vehicle safety routine, slowing vehicle, pulling over to side of road, parking vehicle, and/or idling vehicle, with assistance of ultrasonic sensor(s) **1862**, until emergency vehicle(s) passes.

(268) In at least one embodiment, vehicle **1800** may include CPU(s) **1818** (e.g., discrete CPU(s), or dCPU(s)), that may be coupled to SoC(s) **1804** via a high-speed interconnect (e.g., PCIe). In at least one embodiment, CPU(s) **1818** may include an X86 processor, for example. CPU(s) **1818** may be used to perform any of a variety of functions, including arbitrating potentially inconsistent results between ADAS sensors and SoC(s) **1804**, and/or monitoring status and health of controller(s) **1836** and/or an infotainment system on a chip (“infotainment SoC”) **1830**, for example.

(269) In at least one embodiment, vehicle **1800** may include GPU(s) **1820** (e.g., discrete GPU(s), or dGPU(s)), that may be coupled to SoC(s) **1804** via a high-speed interconnect (e.g., NVIDIA's NVLINK). In at least one embodiment, GPU(s) **1820** may provide additional artificial intelligence functionality, such as by executing redundant and/or different neural networks, and may be used to train and/or update neural networks based at least in part on input (e.g., sensor data) from sensors of vehicle **1800**.

(270) In at least one embodiment, vehicle **1800** may further include network interface **1824** which may include, without limitation, wireless antenna(s) **1826** (e.g., one or more wireless antennas **1826** for different communication protocols, such as a cellular antenna, a Bluetooth antenna, etc.). In at least one embodiment, network interface **1824** may be used to enable wireless connectivity over Internet with cloud (e.g., with server(s) and/or other network devices), with other vehicles, and/or with computing devices (e.g., client devices of passengers). In at least one embodiment, to communicate with other vehicles, a direct link may be established between vehicle **180** and other vehicle and/or an indirect link may be established (e.g., across networks and over Internet). In at least one embodiment, direct links may be provided using a vehicle-to-vehicle communication link. vehicle-to-vehicle communication link may provide vehicle **1800** information about vehicles in proximity to vehicle **1800** (e.g., vehicles in front of, on side of, and/or behind vehicle **1800**). In at least one embodiment, aforementioned functionality may be part of a cooperative adaptive cruise control functionality of vehicle **1800**.

(271) In at least one embodiment, network interface **1824** may include an SoC that provides modulation and demodulation functionality and enables controller(s) **1836** to communicate over wireless networks. In at least one embodiment, network interface **1824** may include a radio frequency front-end for up-conversion from baseband to radio frequency, and down conversion from radio frequency to baseband. In at least one embodiment, frequency conversions may be performed in any technically feasible fashion. For example, frequency conversions could be performed through well-known processes, and/or using super-heterodyne processes. In at least one embodiment, radio frequency front end functionality may be provided by a separate chip. In at least one embodiment, network interface may include wireless functionality for communicating over LTE, WCDMA, UMTS, GSM, CDMA2000, Bluetooth, Bluetooth LE, Wi-Fi, Z-Wave, ZigBee, LoRaWAN, and/or other wireless protocols.

(272) In at least one embodiment, vehicle **1800** may further include data store(s) **1828** which may include, without limitation, off-chip (e.g., off SoC(s) **1804**) storage. In at least one embodiment, data store(s) **1828** may include, without limitation, one or more storage elements including RAM, SRAM, dynamic random-access memory (“DRAM”), video random-access memory (“VRAM”), Flash, hard disks, and/or other components and/or devices that may store at least one bit of data.

(273) In at least one embodiment, vehicle **1800** may further include GNSS sensor(s) **1858** (e.g., GPS and/or assisted GPS sensors), to assist in mapping, perception, occupancy grid generation, and/or path planning functions. In at least one embodiment, any number of GNSS sensor(s) **1858** may be used, including, for example and without limitation, a GPS using a USB connector with an Ethernet to Serial (e.g., RS-232) bridge.

(274) In at least one embodiment, vehicle **1800** may further include RADAR sensor(s) **1860**. RADAR sensor(s) **1860** may be used by vehicle **1800** for long-range vehicle detection, even in darkness and/or severe weather conditions. In at least one embodiment, RADAR functional safety

levels may be ASIL B. RADAR sensor(s) **1860** may use CAN and/or bus **1802** (e.g., to transmit data generated by RADAR sensor(s) **1860**) for control and to access object tracking data, with access to Ethernet to access raw data in some examples. In at least one embodiment, wide variety of RADAR sensor types may be used. For example, and without limitation, RADAR sensor(s) **1860** may be suitable for front, rear, and side RADAR use. In at least one embodiment, one or more of RADAR sensors(s) **1860** are Pulse Doppler RADAR sensor(s).

(275) In at least one embodiment, RADAR sensor(s) **1860** may include different configurations, such as long-range with narrow field of view, short-range with wide field of view, short-range side coverage, etc. In at least one embodiment, long-range RADAR may be used for adaptive cruise control functionality. In at least one embodiment, long-range RADAR systems may provide a broad field of view realized by two or more independent scans, such as within a 250 m range. In at least one embodiment, RADAR sensor(s) **1860** may help in distinguishing between static and moving objects, and may be used by ADAS system **1838** for emergency brake assist and forward collision warning. sensors **1860**(s) included in a long-range RADAR system may include, without limitation, monostatic multimodal RADAR with multiple (e.g., six or more) fixed RADAR antennae and a high-speed CAN and FlexRay interface. In at least one embodiment, with six antennae, central four antennae may create a focused beam pattern, designed to record vehicle's **1800** surroundings at higher speeds with minimal interference from traffic in adjacent lanes. In at least one embodiment, other two antennae may expand field of view, making it possible to quickly detect vehicles entering or leaving vehicle's **1800** lane.

(276) In at least one embodiment, mid-range RADAR systems may include, as an example, a range of up to 160 m (front) or 80 m (rear), and a field of view of up to 42 degrees (front) or 150 degrees (rear). In at least one embodiment, short-range RADAR systems may include, without limitation, any number of RADAR sensor(s) **1860** designed to be installed at both ends of rear bumper. When installed at both ends of rear bumper, in at least one embodiment, a RADAR sensor system may create two beams that constantly monitor blind spot in rear and next to vehicle. In at least one embodiment, short-range RADAR systems may be used in ADAS system **1838** for blind spot detection and/or lane change assist.

(277) In at least one embodiment, vehicle **1800** may further include ultrasonic sensor(s) **1862**. ultrasonic sensor(s) **1862**, which may be positioned at front, back, and/or sides of vehicle **1800**, may be used for park assist and/or to create and update an occupancy grid. In at least one embodiment, a wide variety of ultrasonic sensor(s) **1862** may be used, and different ultrasonic sensor(s) **1862** may be used for different ranges of detection (e.g., 2.5m, 4m). In at least one embodiment, ultrasonic sensor(s) **1862** may operate at functional safety levels of ASIL B.

(278) In at least one embodiment, vehicle **1800** may include LIDAR sensor(s) **1864**. LIDAR sensor(s) **1864** may be used for object and pedestrian detection, emergency braking, collision avoidance, and/or other functions. In at least one embodiment, LIDAR sensor(s) **1864** may be functional safety level ASIL B. In at least one embodiment, vehicle **1800** may include multiple LIDAR sensors **1864** (e.g., two, four, six, etc.) that may use Ethernet (e.g., to provide data to a Gigabit Ethernet switch).

(279) In at least one embodiment, LIDAR sensor(s) **1864** may be capable of providing a list of objects and their distances for a 360-degree field of view. In at least one embodiment, commercially available LIDAR sensor(s) **1864** may have an advertised range of approximately 100 m, with an accuracy of 2 cm-3 cm, and with support for a 100 Mbps Ethernet connection, for example. In at least one embodiment, one or more non-protruding LIDAR sensors **1864** may be used. In such an embodiment, LIDAR sensor(s) **1864** may be implemented as a small device that may be embedded into front, rear, sides, and/or corners of vehicle **1800**. In at least one embodiment, LIDAR sensor(s) **1864**, in such an embodiment, may provide up to a 120-degree horizontal and 35-degree vertical field-of-view, with a 200 m range even for low-reflectivity objects. In at least one embodiment, front-mounted LIDAR sensor(s) **1864** may be configured for a

horizontal field of view between 45 degrees and 135 degrees.

(280) In at least one embodiment, LIDAR technologies, such as 3D flash LIDAR, may also be used. 3D Flash LIDAR uses a flash of a laser as a transmission source, to illuminate surroundings of vehicle **1800** up to approximately 200 m. In at least one embodiment, a flash LIDAR unit includes, without limitation, a receptor, which records laser pulse transit time and reflected light on each pixel, which in turn corresponds to range from vehicle **1800** to objects. In at least one embodiment, flash LIDAR may allow for highly accurate and distortion-free images of surroundings to be generated with every laser flash. In at least one embodiment, four flash LIDAR sensors may be deployed, one at each side of vehicle **1800**. In at least one embodiment, 3D flash LIDAR systems include, without limitation, a solid-state 3D staring array LIDAR camera with no moving parts other than a fan (e.g., a non-scanning LIDAR device). In at least one embodiment, flash LIDAR device may use a 5 nanosecond class I (eye-safe) laser pulse per frame and may capture reflected laser light in form of 3D range point clouds and co-registered intensity data.

(281) In at least one embodiment, vehicle may further include IMU sensor(s) **1866**. In at least one embodiment, IMU sensor(s) **1866** may be located at a center of rear axle of vehicle **1800**, in at least one embodiment. In at least one embodiment, IMU sensor(s) **1866** may include, for example and without limitation, accelerometer(s), magnetometer(s), gyroscope(s), magnetic compass(es), and/or other sensor types. In at least one embodiment, such as in six-axis applications, IMU sensor(s) **1866** may include, without limitation, accelerometers and gyroscopes. In at least one embodiment, such as in nine-axis applications, IMU sensor(s) **1866** may include, without limitation, accelerometers, gyroscopes, and magnetometers.

(282) In at least one embodiment, IMU sensor(s) **1866** may be implemented as a miniature, high performance GPS-Aided Inertial Navigation System (“GPS/INS”) that combines micro-electro-mechanical systems (“MEMS”) inertial sensors, a high-sensitivity GPS receiver, and advanced Kalman filtering algorithms to provide estimates of position, velocity, and attitude. In at least one embodiment, IMU sensor(s) **1866** may enable vehicle **1800** to estimate heading without requiring input from a magnetic sensor by directly observing and correlating changes in velocity from GPS to IMU sensor(s) **1866**. In at least one embodiment, IMU sensor(s) **1866** and GNSS sensor(s) **1858** may be combined in a single integrated unit.

(283) In at least one embodiment, vehicle **1800** may include microphone(s) **1896** placed in and/or around vehicle **1800**. In at least one embodiment, microphone(s) **1896** may be used for emergency vehicle detection and identification, among other things.

(284) In at least one embodiment, vehicle **1800** may further include any number of camera types, including stereo camera(s) **1868**, wide-view camera(s) **1870**, infrared camera(s) **1872**, surround camera(s) **1874**, long-range camera(s) **1898**, mid-range camera(s) **1876**, and/or other camera types. In at least one embodiment, cameras may be used to capture image data around an entire periphery of vehicle **1800**. In at least one embodiment, types of cameras used depends vehicle **1800**. In at least one embodiment, any combination of camera types may be used to provide necessary coverage around vehicle **1800**. In at least one embodiment, number of cameras may differ depending on embodiment. For example, in at least one embodiment, vehicle **1800** could include six cameras, seven cameras, ten cameras, twelve cameras, or another number of cameras. cameras may support, as an example and without limitation, Gigabit Multimedia Serial Link (“GMSL”) and/or Gigabit Ethernet. In at least one embodiment, each of camera(s) is described with more detail previously herein with respect to FIG. **18A** and FIG. **18B**.

(285) In at least one embodiment, vehicle **1800** may further include vibration sensor(s) **1842**. vibration sensor(s) **1842** may measure vibrations of components of vehicle **1800**, such as axle(s). For example, in at least one embodiment, changes in vibrations may indicate a change in road surfaces. In at least one embodiment, when two or more vibration sensors **1842** are used, differences between vibrations may be used to determine friction or slippage of road surface (e.g., when difference in vibration is between a power-driven axle and a freely rotating axle).

(286) In at least one embodiment, vehicle **1800** may include ADAS system **1838**. ADAS system **1838** may include, without limitation, an SoC, in some examples. In at least one embodiment, ADAS system **1838** may include, without limitation, any number and combination of an autonomous/adaptive/automatic cruise control (“ACC”) system, a cooperative adaptive cruise control (“CACC”) system, a forward crash warning (“FCW”) system, an automatic emergency braking (“AEB”) system, a lane departure warning (“LDW”) system, a lane keep assist (“LKA”) system, a blind spot warning (“BSW”) system, a rear cross-traffic warning (“RCTW”) system, a collision warning (“CW”) system, a lane centering (“LC”) system, and/or other systems, features, and/or functionality.

(287) In at least one embodiment, ACC system may use RADAR sensor(s) **1860**, LIDAR sensor(s) **1864**, and/or any number of camera(s). In at least one embodiment, ACC system may include a longitudinal ACC system and/or a lateral ACC system. In at least one embodiment, longitudinal ACC system monitors and controls distance to vehicle immediately ahead of vehicle **1800** and automatically adjust speed of vehicle **1800** to maintain a safe distance from vehicles ahead. In at least one embodiment, lateral ACC system performs distance keeping, and advises vehicle **1800** to change lanes when necessary. In at least one embodiment, lateral ACC is related to other ADAS applications such as LC and CW.

(288) In at least one embodiment, CACC system uses information from other vehicles that may be received via network interface **1824** and/or wireless antenna(s) **1826** from other vehicles via a wireless link, or indirectly, over a network connection (e.g., over Internet). In at least one embodiment, direct links may be provided by a vehicle-to-vehicle (“V2V”) communication link, while indirect links may be provided by an infrastructure-to-vehicle (“I2V”) communication link. In general, V2V communication concept provides information about immediately preceding vehicles (e.g., vehicles immediately ahead of and in same lane as vehicle **1800**), while I2V communication concept provides information about traffic further ahead. In at least one embodiment, CACC system may include either or both I2V and V2V information sources. In at least one embodiment, given information of vehicles ahead of vehicle **1800**, CACC system may be more reliable and it has potential to improve traffic flow smoothness and reduce congestion on road.

(289) In at least one embodiment, FCW system is designed to alert driver to a hazard, so that driver may take corrective action. In at least one embodiment, FCW system uses a front-facing camera and/or RADAR sensor(s) **1860**, coupled to a dedicated processor, DSP, FPGA, and/or ASIC, that is electrically coupled to driver feedback, such as a display, speaker, and/or vibrating component. In at least one embodiment, FCW system may provide a warning, such as in form of a sound, visual warning, vibration and/or a quick brake pulse.

(290) In at least one embodiment, AEB system detects an impending forward collision with another vehicle or other object, and may automatically apply brakes if driver does not take corrective action within a specified time or distance parameter. In at least one embodiment, AEB system may use front-facing camera(s) and/or RADAR sensor(s) **1860**, coupled to a dedicated processor, DSP, FPGA, and/or ASIC. In at least one embodiment, when AEB system detects a hazard, AEB system typically first alerts driver to take corrective action to avoid collision and, if driver does not take corrective action, AEB system may automatically apply brakes in an effort to prevent, or at least mitigate, impact of predicted collision. In at least one embodiment, AEB system, may include techniques such as dynamic brake support and/or crash imminent braking.

(291) In at least one embodiment, LDW system provides visual, audible, and/or tactile warnings, such as steering wheel or seat vibrations, to alert driver when vehicle **1800** crosses lane markings. In at least one embodiment, LDW system does not activate when driver indicates an intentional lane departure, by activating a turn signal. In at least one embodiment, LDW system may use front-side facing cameras, coupled to a dedicated processor, DSP, FPGA, and/or ASIC, that is electrically coupled to driver feedback, such as a display, speaker, and/or vibrating component. In at least one

embodiment, LKA system is a variation of LDW system. LKA system provides steering input or braking to correct vehicle **1800** if vehicle **1800** starts to exit lane.

(292) In at least one embodiment, BSW system detects and warns driver of vehicles in an automobile's blind spot. In at least one embodiment, BSW system may provide a visual, audible, and/or tactile alert to indicate that merging or changing lanes is unsafe. In at least one embodiment, BSW system may provide an additional warning when driver uses a turn signal. In at least one embodiment, BSW system may use rear-side facing camera(s) and/or RADAR sensor(s) **1860**, coupled to a dedicated processor, DSP, FPGA, and/or ASIC, that is electrically coupled to driver feedback, such as a display, speaker, and/or vibrating component.

(293) In at least one embodiment, RCTW system may provide visual, audible, and/or tactile notification when an object is detected outside rear-camera range when vehicle **1800** is backing up. In at least one embodiment, RCTW system includes AEB system to ensure that vehicle brakes are applied to avoid a crash. In at least one embodiment, RCTW system may use one or more rear-facing RADAR sensor(s) **1860**, coupled to a dedicated processor, DSP, FPGA, and/or ASIC, that is electrically coupled to driver feedback, such as a display, speaker, and/or vibrating component.

(294) In at least one embodiment, conventional ADAS systems may be prone to false positive results which may be annoying and distracting to a driver, but typically are not catastrophic, because conventional ADAS systems alert driver and allow driver to decide whether a safety condition truly exists and act accordingly. In at least one embodiment, vehicle **1800** itself decides, in case of conflicting results, whether to heed result from a primary computer or a secondary computer (e.g., first controller **1836** or second controller **1836**). For example, in at least one embodiment, ADAS system **1838** may be a backup and/or secondary computer for providing perception information to a backup computer rationality module. In at least one embodiment, backup computer rationality monitor may run a redundant diverse software on hardware components to detect faults in perception and dynamic driving tasks. In at least one embodiment, outputs from ADAS system **1838** may be provided to a supervisory MCU. In at least one embodiment, if outputs from primary computer and secondary computer conflict, supervisory MCU determines how to reconcile conflict to ensure safe operation.

(295) In at least one embodiment, primary computer may be configured to provide supervisory MCU with a confidence score, indicating primary computer's confidence in chosen result. In at least one embodiment, if confidence score exceeds a threshold, supervisory MCU may follow primary computer's direction, regardless of whether secondary computer provides a conflicting or inconsistent result. In at least one embodiment, where confidence score does not meet threshold, and where primary and secondary computer indicate different results (e.g., a conflict), supervisory MCU may arbitrate between computers to determine appropriate outcome.

(296) In at least one embodiment, supervisory MCU may be configured to run a neural network(s) that is trained and configured to determine, based at least in part on outputs from primary computer and secondary computer, conditions under which secondary computer provides false alarms. In at least one embodiment, neural network(s) in supervisory MCU may learn when secondary computer's output may be trusted, and when it cannot. For example, in at least one embodiment, when secondary computer is a RADAR-based FCW system, a neural network(s) in supervisory MCU may learn when FCW system is identifying metallic objects that are not, in fact, hazards, such as a drainage grate or manhole cover that triggers an alarm. In at least one embodiment, when secondary computer is a camera-based LDW system, a neural network in supervisory MCU may learn to override LDW when bicyclists or pedestrians are present and a lane departure is, in fact, safest maneuver. In at least one embodiment, supervisory MCU may include at least one of a DLA or GPU suitable for running neural network(s) with associated memory. In at least one embodiment, supervisory MCU may comprise and/or be included as a component of SoC(s) **1804**.

(297) In at least one embodiment, ADAS system **1838** may include a secondary computer that performs ADAS functionality using traditional rules of computer vision. In at least one

embodiment, secondary computer may use classic computer vision rules (if-then), and presence of a neural network(s) in supervisory MCU may improve reliability, safety and performance. For example, in at least one embodiment, diverse implementation and intentional non-identity makes overall system more fault-tolerant, especially to faults caused by software (or software-hardware interface) functionality. For example, in at least one embodiment, if there is a software bug or error in software running on primary computer, and non-identical software code running on secondary computer provides same overall result, then supervisory MCU may have greater confidence that overall result is correct, and bug in software or hardware on primary computer is not causing material error.

(298) In at least one embodiment, output of ADAS system **1838** may be fed into primary computer's perception block and/or primary computer's dynamic driving task block. For example, in at least one embodiment, if ADAS system **1838** indicates a forward crash warning due to an object immediately ahead, perception block may use this information when identifying objects. In at least one embodiment, secondary computer may have its own neural network which is trained and thus reduces risk of false positives, as described herein.

(299) In at least one embodiment, vehicle **1800** may further include infotainment SoC **1830** (e.g., an in-vehicle infotainment system (IVI)). Although illustrated and described as an SoC, infotainment system **1830**, in at least one embodiment, may not be an SoC, and may include, without limitation, two or more discrete components. In at least one embodiment, infotainment SoC **1830** may include, without limitation, a combination of hardware and software that may be used to provide audio (e.g., music, a personal digital assistant, navigational instructions, news, radio, etc.), video (e.g., TV, movies, streaming, etc.), phone (e.g., hands-free calling), network connectivity (e.g., LTE, WiFi, etc.), and/or information services (e.g., navigation systems, rear-parking assistance, a radio data system, vehicle related information such as fuel level, total distance covered, brake fuel level, oil level, door open/close, air filter information, etc.) to vehicle **1800**. For example, infotainment SoC **1830** could include radios, disk players, navigation systems, video players, USB and Bluetooth connectivity, carputers, in-car entertainment, WiFi, steering wheel audio controls, hands free voice control, a heads-up display ("HUD"), HMI display **1834**, a telematics device, a control panel (e.g., for controlling and/or interacting with various components, features, and/or systems), and/or other components. In at least one embodiment, infotainment SoC **1830** may further be used to provide information (e.g., visual and/or audible) to user(s) of vehicle, such as information from ADAS system **1838**, autonomous driving information such as planned vehicle maneuvers, trajectories, surrounding environment information (e.g., intersection information, vehicle information, road information, etc.), and/or other information.

(300) In at least one embodiment, infotainment SoC **1830** may include any amount and type of GPU functionality. In at least one embodiment, infotainment SoC **1830** may communicate over bus **1802** (e.g., CAN bus, Ethernet, etc.) with other devices, systems, and/or components of vehicle **1800**. In at least one embodiment, infotainment SoC **1830** may be coupled to a supervisory MCU such that GPU of infotainment system may perform some self-driving functions in event that primary controller(s) **1836** (e.g., primary and/or backup computers of vehicle **1800**) fail. In at least one embodiment, infotainment SoC **1830** may put vehicle **1800** into a chauffeur to safe stop mode, as described herein.

(301) In at least one embodiment, vehicle **1800** may further include instrument cluster **1832** (e.g., a digital dash, an electronic instrument cluster, a digital instrument panel, etc.). instrument cluster **1832** may include, without limitation, a controller and/or supercomputer (e.g., a discrete controller or supercomputer). In at least one embodiment, instrument cluster **1832** may include, without limitation, any number and combination of a set of instrumentation such as a speedometer, fuel level, oil pressure, tachometer, odometer, turn indicators, gearshift position indicator, seat belt warning light(s), parking-brake warning light(s), engine-malfunction light(s), supplemental restraint system (e.g., airbag) information, lighting controls, safety system controls, navigation

information, etc. In some examples, information may be displayed and/or shared among infotainment SoC **1830** and instrument cluster **1832**. In at least one embodiment, instrument cluster **1832** may be included as part of infotainment SoC **1830**, or vice versa.

(302) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in system FIG. **18C** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(303) In at least one embodiment, components, features, and systems of vehicle **1800** in FIG. **18C** include memory storing computer-readable instructions that, if executed, cause one or more processors of components, features, and systems of vehicle **1800** in FIG. **18C** to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. One or more embodiments described elsewhere in this disclosure may be utilized in context of vehicle **1800** of FIG. **18C**, such as techniques described in connection with FIGS. **1-17** and **19-43**.

(304) FIG. **18D** is a diagram of a system **1876** for communication between cloud-based server(s) and autonomous vehicle **1800** of FIG. **18A**, according to at least one embodiment. In at least one embodiment, system **1876** may include, without limitation, server(s) **1878**, network(s) **1890**, and any number and type of vehicles, including vehicle **1800**. server(s) **1878** may include, without limitation, a plurality of GPUs **1884(A)-1884(H)** (collectively referred to herein as GPUs **1884**), PCIe switches **1882(A)-1882(H)** (collectively referred to herein as PCIe switches **1882**), and/or CPUs **1880(A)-1880(B)** (collectively referred to herein as CPUs **1880**). GPUs **1884**, CPUs **1880**, and PCIe switches **1882** may be interconnected with high-speed interconnects such as, for example and without limitation, NVLink interfaces **1888** developed by NVIDIA and/or PCIe connections **1886**. In at least one embodiment, GPUs **1884** are connected via an NVLink and/or NVSwitch SoC and GPUs **1884** and PCIe switches **1882** are connected via PCIe interconnects. In at least one embodiment, although eight GPUs **1884**, two CPUs **1880**, and four PCIe switches **1882** are illustrated, this is not intended to be limiting. In at least one embodiment, each of server(s) **1878** may include, without limitation, any number of GPUs **1884**, CPUs **1880**, and/or PCIe switches **1882**, in any combination. For example, in at least one embodiment, server(s) **1878** could each include eight, sixteen, thirty-two, and/or more GPUs **1884**.

(305) In at least one embodiment, server(s) **1878** may receive, over network(s) **1890** and from vehicles, image data representative of images showing unexpected or changed road conditions, such as recently commenced road-work. In at least one embodiment, server(s) **1878** may transmit, over network(s) **1890** and to vehicles, neural networks **1892**, updated neural networks **1892**, and/or map information **1894**, including, without limitation, information regarding traffic and road conditions. In at least one embodiment, updates to map information **1894** may include, without limitation, updates for HD map **1822**, such as information regarding construction sites, potholes, detours, flooding, and/or other obstructions. In at least one embodiment, neural networks **1892**, updated neural networks **1892**, and/or map information **1894** may have resulted from new training and/or experiences represented in data received from any number of vehicles in environment, and/or based at least in part on training performed at a data center (e.g., using server(s) **1878** and/or other servers).

(306) In at least one embodiment, server(s) **1878** may be used to train machine learning models (e.g., neural networks) based at least in part on training data. training data may be generated by vehicles, and/or may be generated in a simulation (e.g., using a game engine). In at least one embodiment, any amount of training data is tagged (e.g., where associated neural network benefits

from supervised learning) and/or undergoes other pre-processing. In at least one embodiment, any amount of training data is not tagged and/or pre-processed (e.g., where associated neural network does not require supervised learning). In at least one embodiment, once machine learning models are trained, machine learning models may be used by vehicles (e.g., transmitted to vehicles over network(s) **1890**, and/or machine learning models may be used by server(s) **1878** to remotely monitor vehicles.

(307) In at least one embodiment, server(s) **1878** may receive data from vehicles and apply data to up-to-date real-time neural networks for real-time inferencing. In at least one embodiment, server(s) **1878** may include deep-learning supercomputers and/or dedicated AI computers powered by GPU(s) **1884**, such as a DGX and DGX Station machines developed by NVIDIA. However, in at least one embodiment, server(s) **1878** may include deep learning infrastructure that use CPU-powered data centers.

(308) In at least one embodiment, deep-learning infrastructure of server(s) **1878** may be capable of fast, real-time inferencing, and may use that capability to evaluate and verify health of processors, software, and/or associated hardware in vehicle **1800**. For example, in at least one embodiment, deep-learning infrastructure may receive periodic updates from vehicle **1800**, such as a sequence of images and/or objects that vehicle **1800** has located in that sequence of images (e.g., via computer vision and/or other machine learning object classification techniques). In at least one embodiment, deep-learning infrastructure may run its own neural network to identify objects and compare them with objects identified by vehicle **1800** and, if results do not match and deep-learning infrastructure concludes that AI in vehicle **1800** is malfunctioning, then server(s) **1878** may transmit a signal to vehicle **1800** instructing a fail-safe computer of vehicle **1800** to assume control, notify passengers, and complete a safe parking maneuver.

(309) In at least one embodiment, server(s) **1878** may include GPU(s) **1884** and one or more programmable inference accelerators (e.g., NVIDIA's TensorRT **3**). In at least one embodiment, combination of GPU-powered servers and inference acceleration may make real-time responsiveness possible. In at least one embodiment, such as where performance is less critical, servers powered by CPUs, FPGAs, and other processors may be used for inferencing. In at least one embodiment, hardware structure(s) **1515** are used to perform one or more embodiments. Details regarding hardware structure(x) **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**.

(310) Computer Systems

(311) FIG. **19** is a block diagram illustrating an exemplary computer system, which may be a system with interconnected devices and components, a system-on-a-chip (SOC) or some combination thereof **1900** formed with a processor that may include execution units to execute an instruction, according to at least one embodiment. In at least one embodiment, computer system **1900** may include, without limitation, a component, such as a processor **1902** to employ execution units including logic to perform algorithms for process data, in accordance with present disclosure, such as in embodiment described herein. In at least one embodiment, computer system **1900** may include processors, such as PENTIUM® Processor family, Xeon™, Itanium®, XScale™ and/or StrongARM™, Intel® Core™, or Intel® Nervana™ microprocessors available from Intel Corporation of Santa Clara, California, although other systems (including PCs having other microprocessors, engineering workstations, set-top boxes and like) may also be used. In at least one embodiment, computer system **1900** may execute a version of WINDOWS' operating system available from Microsoft Corporation of Redmond, Wash., although other operating systems (UNIX and Linux for example), embedded software, and/or graphical user interfaces, may also be used.

(312) Embodiments may be used in other devices such as handheld devices and embedded applications. Some examples of handheld devices include cellular phones, Internet Protocol devices, digital cameras, personal digital assistants ("PDAs"), and handheld PCs. In at least one

embodiment, embedded applications may include a microcontroller, a digital signal processor (“DSP”), system on a chip, network computers (“NetPCs”), set-top boxes, network hubs, wide area network (“WAN”) switches, or any other system that may perform one or more instructions in accordance with at least one embodiment.

(313) In at least one embodiment, computer system **1900** may include, without limitation, processor **1902** that may include, without limitation, one or more execution units **1908** to perform machine learning model training and/or inferencing according to techniques described herein. In at least one embodiment, system **19** is a single processor desktop or server system, but in another embodiment system **19** may be a multiprocessor system. In at least one embodiment, processor **1902** may include, without limitation, a complex instruction set computer (“CISC”) microprocessor, a reduced instruction set computing (“RISC”) microprocessor, a very long instruction word (“VLIW”) microprocessor, a processor implementing a combination of instruction sets, or any other processor device, such as a digital signal processor, for example. In at least one embodiment, processor **1902** may be coupled to a processor bus **1910** that may transmit data signals between processor **1902** and other components in computer system **1900**.

(314) In at least one embodiment, processor **1902** may include, without limitation, a Level 1 (“L1”) internal cache memory (“cache”) **1904**. In at least one embodiment, processor **1902** may have a single internal cache or multiple levels of internal cache. In at least one embodiment, cache memory may reside external to processor **1902**. Other embodiments may also include a combination of both internal and external caches depending on particular implementation and needs. In at least one embodiment, register file **1906** may store different types of data in various registers including, without limitation, integer registers, floating point registers, status registers, and instruction pointer register.

(315) In at least one embodiment, execution unit **1908**, including, without limitation, logic to perform integer and floating point operations, also resides in processor **1902**. processor **1902** may also include a microcode (“ucode”) read only memory (“ROM”) that stores microcode for certain macro instructions. In at least one embodiment, execution unit **1908** may include logic to handle a packed instruction set **1909**. In at least one embodiment, by including packed instruction set **1909** in instruction set of a general-purpose processor **1902**, along with associated circuitry to execute instructions, operations used by many multimedia applications may be performed using packed data in a general-purpose processor **1902**. In one or more embodiments, many multimedia applications may be accelerated and executed more efficiently by using full width of a processor's data bus for performing operations on packed data, which may eliminate need to transfer smaller units of data across processor's data bus to perform one or more operations one data element at a time.

(316) In at least one embodiment, execution unit **1908** may also be used in microcontrollers, embedded processors, graphics devices, DSPs, and other types of logic circuits. In at least one embodiment, computer system **1900** may include, without limitation, a memory **1920**. In at least one embodiment, memory **1920** may be implemented as a Dynamic Random Access Memory (“DRAM”) device, a Static Random Access Memory (“SRAM”) device, flash memory device, or other memory device. memory **1920** may store instruction(s) **1919** and/or data **1921** represented by data signals that may be executed by processor **1902**.

(317) In at least one embodiment, system logic chip may be coupled to processor bus **1910** and memory **1920**. In at least one embodiment, system logic chip may include, without limitation, a memory controller hub (“MCH”) **1916**, and processor **1902** may communicate with MCH **1916** via processor bus **1910**. In at least one embodiment, MCH **1916** may provide a high bandwidth memory path **1918** to memory **1920** for instruction and data storage and for storage of graphics commands, data and textures. In at least one embodiment, MCH **1916** may direct data signals between processor **1902**, memory **1920**, and other components in computer system **1900** and to bridge data signals between processor bus **1910**, memory **1920**, and a system I/O **1922**. In at least

one embodiment, system logic chip may provide a graphics port for coupling to a graphics controller. In at least one embodiment, MCH **1916** may be coupled to memory **1920** through a high bandwidth memory path **1918** and graphics/video card **1912** may be coupled to MCH **1916** through an Accelerated Graphics Port (“AGP”) interconnect **1914**.

(318) In at least one embodiment, computer system **1900** may use system I/O **1922** that is a proprietary hub interface bus to couple MCH **1916** to I/O controller hub (“ICH”) **1930**. In at least one embodiment, ICH **1930** may provide direct connections to some I/O devices via a local I/O bus. In at least one embodiment, local I/O bus may include, without limitation, a high-speed I/O bus for connecting peripherals to memory **1920**, chipset, and processor **1902**. Examples may include, without limitation, an audio controller **1929**, a firmware hub (“flash BIOS”) **1928**, a wireless transceiver **1926**, a data storage **1924**, a legacy I/O controller **1923** containing user input and keyboard interfaces, a serial expansion port **1927**, such as Universal Serial Bus (“USB”), and a network controller **1934**. data storage **1924** may comprise a hard disk drive, a floppy disk drive, a CD-ROM device, a flash memory device, or other mass storage device.

(319) In at least one embodiment, FIG. **19** illustrates a system, which includes interconnected hardware devices or “chips”, whereas in other embodiments, FIG. **19** may illustrate an exemplary System on a Chip (“SoC”). In at least one embodiment, devices illustrated in FIG. **19** may be interconnected with proprietary interconnects, standardized interconnects (e.g., PCIe) or some combination thereof. In at least one embodiment, one or more components of system **1900** are interconnected using compute express link (CXL) interconnects.

(320) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in system FIG. **19** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(321) In at least one embodiment, computer system **1900** includes memory storing computer-readable executable instruction that, as a result of execution, causes one or more processors to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, computer system **1900** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Computer system **1900** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-18** and **20-43**.

(322) FIG. **20** is a block diagram illustrating an electronic device **2000** for utilizing a processor **2010**, according to at least one embodiment. In at least one embodiment, electronic device **2000** may be, for example and without limitation, a notebook, a tower server, a rack server, a blade server, a laptop, a desktop, a tablet, a mobile device, a phone, an embedded computer, or any other suitable electronic device.

(323) In at least one embodiment, system **2000** may include, without limitation, processor **2010** communicatively coupled to any suitable number or kind of components, peripherals, modules, or devices. In at least one embodiment, processor **2010** coupled using a bus or interface, such as a 1st C. bus, a System Management Bus (“SMBus”), a Low Pin Count (LPC) bus, a Serial Peripheral Interface (“SPI”), a High Definition Audio (“HDA”) bus, a Serial Advance Technology Attachment (“SATA”) bus, a Universal Serial Bus (“USB”) (versions 1, 2, 3), or a Universal Asynchronous Receiver/Transmitter (“UART”) bus. In at least one embodiment, FIG. **20** illustrates a system, which includes interconnected hardware devices or “chips”, whereas in other embodiments, FIG.

20 may illustrate an exemplary System on a Chip (“SoC”). In at least one embodiment, devices illustrated in FIG. **20** may be interconnected with proprietary interconnects, standardized interconnects (e.g., PCIe) or some combination thereof. In at least one embodiment, one or more components of FIG. **20** are interconnected using compute express link (CXL) interconnects.

(324) In at least one embodiment, FIG. **20** may include a display **2024**, a touch screen **2025**, a touch pad **2030**, a Near Field Communications unit (“NFC”) **2045**, a sensor hub **2040**, a thermal sensor **2046**, an Express Chipset (“EC”) **2035**, a Trusted Platform Module (“TPM”) **2038**, BIOS/firmware/flash memory (“BIOS, FW Flash”) **2022**, a DSP **2060**, a drive “SSD or HDD”) **2020** such as a Solid State Disk (“SSD”) or a Hard Disk Drive (“HDD”), a wireless local area network unit (“WLAN”) **2050**, a Bluetooth unit **2052**, a Wireless Wide Area Network unit (“WWAN”) **2056**, a Global Positioning System (GPS) **2055**, a camera (“USB 3.0 camera”) **2054** such as a USB 3.0 camera, or a Low Power Double Data Rate (“LPDDR”) memory unit (“LPDDR3”) **2015** implemented in, for example, LPDDR3 standard. These components may each be implemented in any suitable manner.

(325) In at least one embodiment, other components may be communicatively coupled to processor **2010** through components discussed above. In at least one embodiment, an accelerometer **2041**, Ambient Light Sensor (“ALS”) **2042**, compass **2043**, and a gyroscope **2044** may be communicatively coupled to sensor hub **2040**. In at least one embodiment, thermal sensor **2039**, a fan **2037**, a keyboard **2046**, and a touch pad **2030** may be communicatively coupled to EC **2035**. In at least one embodiment, speaker **2063**, a headphones **2064**, and a microphone (“mic”) **2065** may be communicatively coupled to an audio unit (“audio codec and class d amp”) **2064**, which may in turn be communicatively coupled to DSP **2060**. In at least one embodiment, audio unit **2064** may include, for example and without limitation, an audio coder/decoder (“codec”) and a class D amplifier. In at least one embodiment, SIM card (“SIM”) **2057** may be communicatively coupled to WWAN unit **2056**. In at least one embodiment, components such as WLAN unit **2050** and Bluetooth unit **2052**, as well as WWAN unit **2056** may be implemented in a Next Generation Form Factor (“NGFF”).

(326) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in system FIG. **20** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(327) In at least one embodiment, electronic device **2000** includes memory storing computer-readable executable instruction that, as a result of execution, causes one or more processors to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, electronic device **2000** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Electronic device **2000** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-19** and **21-43**.

(328) FIG. **21** illustrates a computer system **2100**, according to at least one embodiment. In at least one embodiment, computer system **2100** is configured to implement various processes and methods described throughout this disclosure.

(329) In at least one embodiment, computer system **2100** comprises, without limitation, at least one central processing unit (“CPU”) **2102** that is connected to a communication bus **2110** implemented using any suitable protocol, such as PCI (“Peripheral Component Interconnect”), peripheral

component interconnect express (“PCI-Express”), AGP (“Accelerated Graphics Port”), HyperTransport, or any other bus or point-to-point communication protocol(s). In at least one embodiment, computer system **2100** includes, without limitation, a main memory **2104** and control logic (e.g., implemented as hardware, software, or a combination thereof) and data are stored in main memory **2104** which may take form of random access memory (“RAM”). In at least one embodiment, a network interface subsystem (“network interface”) **2122** provides an interface to other computing devices and networks for receiving data from and transmitting data to other systems from computer system **2100**.

(330) In at least one embodiment, computer system **2100**, in at least one embodiment, includes, without limitation, input devices **2108**, parallel processing system **2112**, and display devices **2106** which can be implemented using a conventional cathode ray tube (“CRT”), liquid crystal display (“LCD”), light emitting diode (“LED”), plasma display, or other suitable display technologies. In at least one embodiment, user input is received from input devices **2108** such as keyboard, mouse, touchpad, microphone, and more. In at least one embodiment, each of foregoing modules can be situated on a single semiconductor platform to form a processing system.

(331) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in system FIG. **21** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(332) In at least one embodiment, computer system **2100** includes memory storing computer-readable executable instruction that, as a result of execution, causes one or more processors to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, computer system **2100** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Computer system **2100** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-20** and **22-43**.

(333) FIG. **22** illustrates a computer system **2200**, according to at least one embodiment. In at least one embodiment, computer system **2200** includes, without limitation, a computer **2210** and a USB stick **2220**. In at least one embodiment, computer **2210** may include, without limitation, any number and type of processor(s) (not shown) and a memory (not shown). In at least one embodiment, computer **2210** includes, without limitation, a server, a cloud instance, a laptop, and a desktop computer.

(334) In at least one embodiment, USB stick **2220** includes, without limitation, a processing unit **2230**, a USB interface **2240**, and USB interface logic **2250**. In at least one embodiment, processing unit **2230** may be any instruction execution system, apparatus, or device capable of executing instructions. In at least one embodiment, processing unit **2230** may include, without limitation, any number and type of processing cores (not shown). In at least one embodiment, processing core **2230** comprises an application specific integrated circuit (“ASIC”) that is optimized to perform any amount and type of operations associated with machine learning. For instance, in at least one embodiment, processing core **2230** is a tensor processing unit (“TPC”) that is optimized to perform machine learning inference operations. In at least one embodiment, processing core **2230** is a vision processing unit (“VPU”) that is optimized to perform machine vision and machine learning inference operations.

(335) In at least one embodiment, USB interface **2240** may be any type of USB connector or USB

socket. For instance, in at least one embodiment, USB interface **2240** is a USB 3.0 Type-C socket for data and power. In at least one embodiment, USB interface **2240** is a USB 3.0 Type-A connector. In at least one embodiment, USB interface logic **2250** may include any amount and type of logic that enables processing unit **2230** to interface with or devices (e.g., computer **2210**) via USB connector **2240**.

(336) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in system FIG. **22** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(337) In at least one embodiment, computer system **2200** includes memory storing computer-readable executable instruction that, as a result of execution, causes one or more processors to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, computer system **2200** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Computer system **2200** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-21** and **23-43**.

(338) FIG. **23A** illustrates an exemplary architecture in which a plurality of GPUs **2310-2313** is communicatively coupled to a plurality of multi-core processors **2305-2306** over high-speed links **2340-2343** (e.g., buses, point-to-point interconnects, etc.). In one embodiment, high-speed links **2340-2343** support a communication throughput of 4 GB/s, 30 GB/s, 80 GB/s or higher. Various interconnect protocols may be used including, but not limited to, PCIe 4.0 or 5.0 and NVLink 2.0.

(339) In addition, and in one embodiment, two or more of GPUs **2310-2313** are interconnected over high-speed links **2329-2330**, which may be implemented using same or different protocols/links than those used for high-speed links **2340-2343**. Similarly, two or more of multi-core processors **2305-2306** may be connected over high speed link **2328** which may be symmetric multi-processor (SMP) buses operating at 20 GB/s, 30 GB/s, 120 GB/s or higher. Alternatively, all communication between various system components shown in FIG. **23A** may be accomplished using same protocols/links (e.g., over a common interconnection fabric).

(340) In one embodiment, each multi-core processor **2305-2306** is communicatively coupled to a processor memory **2301-2302**, via memory interconnects **2326-2327**, respectively, and each GPU **2310-2313** is communicatively coupled to GPU memory **2320-2323** over GPU memory interconnects **2350-2353**, respectively. Memory interconnects **2326-2327** and **2350-2353** may utilize same or different memory access technologies. By way of example, and not limitation, processor memories **2301-2302** and GPU memories **2320-2323** may be volatile memories such as dynamic random access memories (DRAMs) (including stacked DRAMs), Graphics DDR SDRAM (GDDR) (e.g., GDDR5, GDDR6), or High Bandwidth Memory (HBM) and/or may be non-volatile memories such as 3D XPoint or Nano-Ram. In one embodiment, some portion of processor memories **2301-2302** may be volatile memory and another portion may be non-volatile memory (e.g., using a two-level memory (2LM) hierarchy).

(341) As described herein, although various processors **2305-2306** and GPUs **2310-2313** may be physically coupled to a particular memory **2301-2302**, **2320-2323**, respectively, a unified memory architecture may be implemented in which a same virtual system address space (also referred to as “effective address” space) is distributed among various physical memories. For example, processor memories **2301-2302** may each comprise 64 GB of system memory address space and GPU

memories **2320-2323** may each comprise 32 GB of system memory address space (resulting in a total of 256 GB addressable memory in this example).

(342) FIG. **23B** illustrates additional details for an interconnection between a multi-core processor **2307** and a graphics acceleration module **2346** in accordance with one exemplary embodiment. Graphics acceleration module **2346** may include one or more GPU chips integrated on a line card which is coupled to processor **2307** via high-speed link **2340**. Alternatively, graphics acceleration module **2346** may be integrated on a same package or chip as processor **2307**.

(343) In at least one embodiment, illustrated processor **2307** includes a plurality of cores **2360A-2360D**, each with a translation lookaside buffer **2361A-2361D** and one or more caches **2362A-2362D**. In at least one embodiment, cores **2360A-2360D** may include various other components for executing instructions and processing data which are not illustrated. Caches **2362A-2362D** may comprise level 1 (L1) and level 2 (L2) caches. In addition, one or more shared caches **2356** may be included in caches **2362A-2362D** and shared by sets of cores **2360A-2360D**. For example, one embodiment of processor **2307** includes 24 cores, each with its own L1 cache, twelve shared L2 caches, and twelve shared L3 caches. In this embodiment, one or more L2 and L3 caches are shared by two adjacent cores. Processor **2307** and graphics acceleration module **2346** connect with system memory **2314**, which may include processor memories **2301-2302** of FIG. **23A**.

(344) Coherency is maintained for data and instructions stored in various caches **2362A-2362D**, **2356** and system memory **2314** via inter-core communication over a coherence bus **2364**. For example, each cache may have cache coherency logic/circuitry associated therewith to communicate to over coherence bus **2364** in response to detected reads or writes to particular cache lines. In one implementation, a cache snooping protocol is implemented over coherence bus **2364** to snoop cache accesses.

(345) In one embodiment, a proxy circuit **2325** communicatively couples graphics acceleration module **2346** to coherence bus **2364**, allowing graphics acceleration module **2346** to participate in a cache coherence protocol as a peer of cores **2360A-2360D**. In particular, an interface **2335** provides connectivity to proxy circuit **2325** over high-speed link **2340** (e.g., a PCIe bus, NVLink, etc.) and an interface **2337** connects graphics acceleration module **2346** to link **2340**.

(346) In one implementation, an accelerator integration circuit **2336** provides cache management, memory access, context management, and interrupt management services on behalf of a plurality of graphics processing engines **2331, 2332, N** of graphics acceleration module **2346**. Graphics processing engines **2331, 2332, N** may each comprise a separate graphics processing unit (GPU). Alternatively, graphics processing engines **2331, 2332, N** may comprise different types of graphics processing engines within a GPU such as graphics execution units, media processing engines (e.g., video encoders/decoders), samplers, and blit engines. In at least one embodiment, graphics acceleration module **2346** may be a GPU with a plurality of graphics processing engines **2331-2332, N** or graphics processing engines **2331-2332, N** may be individual GPUs integrated on a common package, line card, or chip.

(347) In one embodiment, accelerator integration circuit **2336** includes a memory management unit (MMU) **2339** for performing various memory management functions such as virtual-to-physical memory translations (also referred to as effective-to-real memory translations) and memory access protocols for accessing system memory **2314**. MMU **2339** may also include a translation lookaside buffer (TLB) (not shown) for caching virtual/effective to physical/real address translations. In one implementation, a cache **2338** stores commands and data for efficient access by graphics processing engines **2331-2332, N**. In one embodiment, data stored in cache **2338** and graphics memories **2333-2334, M** is kept coherent with core caches **2362A-2362D, 2356** and system memory **2314**. As mentioned, this may be accomplished via proxy circuit **2325** on behalf of cache **2338** and memories **2333-2334, M** (e.g., sending updates to cache **2338** related to modifications/accesses of cache lines on processor caches **2362A-2362D, 2356** and receiving updates from cache **2338**).

(348) A set of registers **2345** store context data for threads executed by graphics processing engines **2331-2332**, N and a context management circuit **2348** manages thread contexts. For example, context management circuit **2348** may perform save and restore operations to save and restore contexts of various threads during contexts switches (e.g., where a first thread is saved and a second thread is stored so that a second thread can be execute by a graphics processing engine). For example, on a context switch, context management circuit **2348** may store current register values to a designated region in memory (e.g., identified by a context pointer). It may then restore register values when returning to a context. In one embodiment, an interrupt management circuit **2347** receives and processes interrupts received from system devices.

(349) In one implementation, virtual/effective addresses from a graphics processing engine **2331** are translated to real/physical addresses in system memory **2314** by MMU **2339**. One embodiment of accelerator integration circuit **2336** supports multiple (e.g., 4, 8, 16) graphics accelerator modules **2346** and/or other accelerator devices. Graphics accelerator module **2346** may be dedicated to a single application executed on processor **2307** or may be shared between multiple applications. In one embodiment, a virtualized graphics execution environment is presented in which resources of graphics processing engines **2331-2332**, N are shared with multiple applications or virtual machines (VMs). In at least one embodiment, resources may be subdivided into “slices” which are allocated to different VMs and/or applications based on processing requirements and priorities associated with VMs and/or applications.

(350) In at least one embodiment, accelerator integration circuit **2336** performs as a bridge to a system for graphics acceleration module **2346** and provides address translation and system memory cache services. In addition, accelerator integration circuit **2336** may provide virtualization facilities for a host processor to manage virtualization of graphics processing engines **2331-2332**, interrupts, and memory management.

(351) Because hardware resources of graphics processing engines **2331-2332**, N are mapped explicitly to a real address space seen by host processor **2307**, any host processor can address these resources directly using an effective address value. One function of accelerator integration circuit **2336**, in one embodiment, is physical separation of graphics processing engines **2331-2332**, N so that they appear to a system as independent units.

(352) In at least one embodiment, one or more graphics memories **2333-2334**, M are coupled to each of graphics processing engines **2331-2332**, N, respectively. Graphics memories **2333-2334**, M store instructions and data being processed by each of graphics processing engines **2331-2332**, N. Graphics memories **2333-2334**, M may be volatile memories such as DRAMs (including stacked DRAMs), GDDR memory (e.g., GDDR5, GDDR6), or HBM, and/or may be non-volatile memories such as 3D XPoint or Nano-Ram.

(353) In one embodiment, to reduce data traffic over link **2340**, biasing techniques are used to ensure that data stored in graphics memories **2333-2334**, M is data which will be used most frequently by graphics processing engines **2331-2332**, N and preferably not used by cores **2360A-2360D** (at least not frequently). Similarly, a biasing mechanism attempts to keep data needed by cores (and preferably not graphics processing engines **2331-2332**, N) within caches **2362A-2362D**, **2356** of cores and system memory **2314**.

(354) FIG. **23C** illustrates another exemplary embodiment in which accelerator integration circuit **2336** is integrated within processor **2307**. In this embodiment, graphics processing engines **2331-2332**, N communicate directly over high-speed link **2340** to accelerator integration circuit **2336** via interface **2337** and interface **2335** (which, again, may be utilize any form of bus or interface protocol). Accelerator integration circuit **2336** may perform same operations as those described with respect to FIG. **23B**, but potentially at a higher throughput given its close proximity to coherence bus **2364** and caches **2362A-2362D**, **2356**. One embodiment supports different programming models including a dedicated-process programming model (no graphics acceleration module virtualization) and shared programming models (with virtualization), which may include

programming models which are controlled by accelerator integration circuit **2336** and programming models which are controlled by graphics acceleration module **2346**.

(355) In at least one embodiment, graphics processing engines **2331-2332**, N are dedicated to a single application or process under a single operating system. In at least one embodiment, a single application can funnel other application requests to graphics processing engines **2331-2332**, N, providing virtualization within a VM/partition.

(356) In at least one embodiment, graphics processing engines **2331-2332**, N, may be shared by multiple VM/application partitions. In at least one embodiment, shared models may use a system hypervisor to virtualize graphics processing engines **2331-2332**, N to allow access by each operating system. For single-partition systems without a hypervisor, graphics processing engines **2331-2332**, N are owned by an operating system. In at least one embodiment, an operating system can virtualize graphics processing engines **2331-2332**, N to provide access to each process or application.

(357) In at least one embodiment, graphics acceleration module **2346** or an individual graphics processing engine **2331-2332**, N selects a process element using a process handle. In one embodiment, process elements are stored in system memory **2314** and are addressable using an effective address to real address translation techniques described herein. In at least one embodiment, a process handle may be an implementation-specific value provided to a host process when registering its context with graphics processing engine **2331-2332**, N (that is, calling system software to add a process element to a process element linked list). In at least one embodiment, a lower 16-bits of a process handle may be an offset of the process element within a process element linked list.

(358) FIG. **23D** illustrates an exemplary accelerator integration slice **2390**. As used herein, a “slice” comprises a specified portion of processing resources of accelerator integration circuit **2336**.

Application effective address space **2382** within system memory **2314** stores process elements **2383**. In one embodiment, process elements **2383** are stored in response to GPU invocations **2381** from applications **2380** executed on processor **2307**. A process element **2383** contains process state for corresponding application **2380**. A work descriptor (WD) **2384** contained in process element **2383** can be a single job requested by an application or may contain a pointer to a queue of jobs. In at least one embodiment, WD **2384** is a pointer to a job request queue in an application's address space **2382**.

(359) Graphics acceleration module **2346** and/or individual graphics processing engines **2331-2332**, N can be shared by all or a subset of processes in a system. In at least one embodiment, an infrastructure for setting up process state and sending a WD **2384** to a graphics acceleration module **2346** to start a job in a virtualized environment may be included.

(360) In at least one embodiment, a dedicated-process programming model is implementation-specific. In this model, a single process owns graphics acceleration module **2346** or an individual graphics processing engine **2331**. Because graphics acceleration module **2346** is owned by a single process, a hypervisor initializes accelerator integration circuit **2336** for an owning partition and an operating system initializes accelerator integration circuit **2336** for an owning process when graphics acceleration module **2346** is assigned.

(361) In operation, a WD fetch unit **2391** in accelerator integration slice **2390** fetches next WD **2384** which includes an indication of work to be done by one or more graphics processing engines of graphics acceleration module **2346**. Data from WD **2384** may be stored in registers **2345** and used by MMU **2339**, interrupt management circuit **2347** and/or context management circuit **2348** as illustrated. For example, one embodiment of MMU **2339** includes segment/page walk circuitry for accessing segment/page tables **2386** within OS virtual address space **2385**. Interrupt management circuit **2347** may process interrupt events **2392** received from graphics acceleration module **2346**. When performing graphics operations, an effective address **2393** generated by a graphics processing engine **2331-2332**, N is translated to a real address by MMU **2339**.

(362) In one embodiment, a same set of registers **2345** are duplicated for each graphics processing engine **2331-2332**, N and/or graphics acceleration module **2346** and may be initialized by a hypervisor or operating system. Each of these duplicated registers may be included in an accelerator integration slice **2390**. Exemplary registers that may be initialized by a hypervisor are shown in Table 1.

(363) TABLE-US-00015 TABLE 1 Hypervisor Initialized Registers 1 Slice Control Register 2 Real Address (RA) Scheduled Processes Area Pointer 3 Authority Mask Override Register 4 Interrupt Vector Table Entry Offset 5 Interrupt Vector Table Entry Limit 6 State Register 7 Logical Partition ID 8 Real address (RA) Hypervisor Accelerator Utilization Record Pointer 9 Storage Description Register

(364) Exemplary registers that may be initialized by an operating system are shown in Table 2.

(365) TABLE-US-00016 TABLE 2 Operating System Initialized Registers 1 Process and Thread Identification 2 Effective Address (EA) Context Save/Restore Pointer 3 Virtual Address (VA) Accelerator Utilization Record Pointer 4 Virtual Address (VA) Storage Segment Table Pointer 5 Authority Mask 6 Work descriptor

(366) In one embodiment, each WD **2384** is specific to a particular graphics acceleration module **2346** and/or graphics processing engines **2331-2332**, N. It contains all information required by a graphics processing engine **2331-2332**, N to do work or it can be a pointer to a memory location where an application has set up a command queue of work to be completed.

(367) FIG. **23E** illustrates additional details for one exemplary embodiment of a shared model. This embodiment includes a hypervisor real address space **2398** in which a process element list **2399** is stored. Hypervisor real address space **2398** is accessible via a hypervisor **2396** which virtualizes graphics acceleration module engines for operating system **2395**.

(368) In at least one embodiment, shared programming models allow for all or a subset of processes from all or a subset of partitions in a system to use a graphics acceleration module **2346**. There are two programming models where graphics acceleration module **2346** is shared by multiple processes and partitions: time-sliced shared and graphics directed shared.

(369) In this model, system hypervisor **2396** owns graphics acceleration module **2346** and makes its function available to all operating systems **2395**. For a graphics acceleration module **2346** to support virtualization by system hypervisor **2396**, graphics acceleration module **2346** may adhere to the following: 1) An application's job request must be autonomous (that is, state does not need to be maintained between jobs), or graphics acceleration module **2346** must provide a context save and restore mechanism. 2) An application's job request is guaranteed by graphics acceleration module **2346** to complete in a specified amount of time, including any translation faults, or graphics acceleration module **2346** provides an ability to preempt processing of a job. 3) Graphics acceleration module **2346** must be guaranteed fairness between processes when operating in a directed shared programming model.

(370) In at least one embodiment, application **2380** is required to make an operating system **2395** system call with a graphics acceleration module **2346** type, a work descriptor (WD), an authority mask register (AMR) value, and a context save/restore area pointer (CSRP). In at least one embodiment, graphics acceleration module **2346** type describes a targeted acceleration function for a system call. In at least one embodiment, graphics acceleration module **2346** type may be a system-specific value. In at least one embodiment, WD is formatted specifically for graphics acceleration module **2346** and can be in a form of a graphics acceleration module **2346** command, an effective address pointer to a user-defined structure, an effective address pointer to a queue of commands, or any other data structure to describe work to be done by graphics acceleration module **2346**. In one embodiment, an AMR value is an AMR state to use for a current process. In at least one embodiment, a value passed to an operating system is similar to an application setting an AMR. If accelerator integration circuit **2336** and graphics acceleration module **2346** implementations do not support a User Authority Mask Override Register (UAMOR), an operating

system may apply a current UAMOR value to an AMR value before passing an AMR in a hypervisor call. Hypervisor **2396** may optionally apply a current Authority Mask Override Register (AMOR) value before placing an AMR into process element **2383**. In at least one embodiment, CSRP is one of registers **2345** containing an effective address of an area in an application's address space **2382** for graphics acceleration module **2346** to save and restore context state. This pointer is optional if no state is required to be saved between jobs or when a job is preempted. In at least one embodiment, context save/restore area may be pinned system memory.

(371) Upon receiving a system call, operating system **2395** may verify that application **2380** has registered and been given authority to use graphics acceleration module **2346**. Operating system **2395** then calls hypervisor **2396** with information shown in Table 3.

(372) TABLE-US-00017 TABLE 3 OS to Hypervisor Call Parameters 1 A work descriptor (WD) 2 An Authority Mask Register (AMR) value (potentially masked) 3 An effective address (EA) Context Save/Restore Area Pointer (CSRP) 4 A process ID (PID) and optional thread ID (TID) 5 A virtual address (VA) accelerator utilization record pointer (AURP) 6 Virtual address of storage segment table pointer (SSTP) 7 A logical interrupt service number (LISN)

(373) Upon receiving a hypervisor call, hypervisor **2396** verifies that operating system **2395** has registered and been given authority to use graphics acceleration module **2346**. Hypervisor **2396** then puts process element **2383** into a process element linked list for a corresponding graphics acceleration module **2346** type. A process element may include information shown in Table 4.

(374) TABLE-US-00018 TABLE 4 Process Element Information 1 A work descriptor (WD) 2 An Authority Mask Register (AMR) value (potentially masked). 3 An effective address (EA) Context Save/Restore Area Pointer (CSRP) 4 A process ID (PID) and optional thread ID (TID) 5 A virtual address (VA) accelerator utilization record pointer (AURP) 6 Virtual address of storage segment table pointer (SSTP) 7 A logical interrupt service number (LISN) 8 Interrupt vector table, derived from hypervisor call parameters 9 A state register (SR) value 10 A logical partition ID (LPID) 11 A real address (RA) hypervisor accelerator utilization record pointer 12 Storage Descriptor Register (SDR)

(375) In at least one embodiment, hypervisor initializes a plurality of accelerator integration slice **2390** registers **2345**.

(376) As illustrated in FIG. **23F**, in at least one embodiment, a unified memory is used, addressable via a common virtual memory address space used to access physical processor memories **2301-2302** and GPU memories **2320-2323**. In this implementation, operations executed on GPUs **2310-2313** utilize a same virtual/effective memory address space to access processor memories **2301-2302** and vice versa, thereby simplifying programmability. In one embodiment, a first portion of a virtual/effective address space is allocated to processor memory **2301**, a second portion to second processor memory **2302**, a third portion to GPU memory **2320**, and so on. In at least one embodiment, an entire virtual/effective memory space (sometimes referred to as an effective address space) is thereby distributed across each of processor memories **2301-2302** and GPU memories **2320-2323**, allowing any processor or GPU to access any physical memory with a virtual address mapped to that memory.

(377) In one embodiment, bias/coherence management circuitry **2394A-2394E** within one or more of MMUs **2339A-2339E** ensures cache coherence between caches of one or more host processors (e.g., **2305**) and GPUs **2310-2313** and implements biasing techniques indicating physical memories in which certain types of data should be stored. While multiple instances of bias/coherence management circuitry **2394A-2394E** are illustrated in FIG. **23F**, bias/coherence circuitry may be implemented within an MMU of one or more host processors **2305** and/or within accelerator integration circuit **2336**.

(378) One embodiment allows GPU-attached memory **2320-2323** to be mapped as part of system memory, and accessed using shared virtual memory (SVM) technology, but without suffering performance drawbacks associated with full system cache coherence. In at least one embodiment,

an ability for GPU-attached memory **2320-2323** to be accessed as system memory without onerous cache coherence overhead provides a beneficial operating environment for GPU offload. This arrangement allows host processor **2305** software to setup operands and access computation results, without overhead of tradition I/O DMA data copies. Such traditional copies involve driver calls, interrupts and memory mapped I/O (MMIO) accesses that are all inefficient relative to simple memory accesses. In at least one embodiment, an ability to access GPU attached memory **2320-2323** without cache coherence overheads can be critical to execution time of an offloaded computation. In cases with substantial streaming write memory traffic, for example, cache coherence overhead can significantly reduce an effective write bandwidth seen by a GPU **2310-2313**. In at least one embodiment, efficiency of operand setup, efficiency of results access, and efficiency of GPU computation may play a role in determining effectiveness of a GPU offload. (379) In at least one embodiment, selection of GPU bias and host processor bias is driven by a bias tracker data structure. A bias table may be used, for example, which may be a page-granular structure (i.e., controlled at a granularity of a memory page) that includes 1 or 2 bits per GPU-attached memory page. In at least one embodiment, a bias table may be implemented in a stolen memory range of one or more GPU-attached memories **2320-2323**, with or without a bias cache in GPU **2310-2313** (e.g., to cache frequently/recently used entries of a bias table). Alternatively, an entire bias table may be maintained within a GPU.

(380) In at least one embodiment, a bias table entry associated with each access to GPU-attached memory **2320-2323** is accessed prior to actual access to a GPU memory, causing the following operations. First, local requests from GPU **2310-2313** that find their page in GPU bias are forwarded directly to a corresponding GPU memory **2320-2323**. Local requests from a GPU that find their page in host bias are forwarded to processor **2305** (e.g., over a high-speed link as discussed above). In one embodiment, requests from processor **2305** that find a requested page in host processor bias complete a request like a normal memory read. Alternatively, requests directed to a GPU-biased page may be forwarded to GPU **2310-2313**. In at least one embodiment, a GPU may then transition a page to a host processor bias if it is not currently using a page. In at least one embodiment, bias state of a page can be changed either by a software-based mechanism, a hardware-assisted software-based mechanism, or, for a limited set of cases, a purely hardware-based mechanism.

(381) One mechanism for changing bias state employs an API call (e.g. OpenCL), which, in turn, calls a GPU's device driver which, in turn, sends a message (or enqueues a command descriptor) to a GPU directing it to change a bias state and, for some transitions, perform a cache flushing operation in a host. In at least one embodiment, cache flushing operation is used for a transition from host processor **2305** bias to GPU bias, but is not for an opposite transition.

(382) In one embodiment, cache coherency is maintained by temporarily rendering GPU-biased pages uncacheable by host processor **2305**. To access these pages, processor **2305** may request access from GPU **2310** which may or may not grant access right away. Thus, to reduce communication between processor **2305** and GPU **2310** it is beneficial to ensure that GPU-biased pages are those which are required by a GPU but not host processor **2305** and vice versa.

(383) Hardware structure(s) **1515** are used to perform one or more embodiments. Details regarding the hardware structure(x) **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**.

(384) FIG. **24** illustrates exemplary integrated circuits and associated graphics processors that may be fabricated using one or more IP cores, according to various embodiments described herein. In addition to what is illustrated, other logic and circuits may be included in at least one embodiment, including additional graphics processors/cores, peripheral interface controllers, or general-purpose processor cores.

(385) FIG. **24** is a block diagram illustrating an exemplary system on a chip integrated circuit **2400** that may be fabricated using one or more IP cores, according to at least one embodiment. In at least one embodiment, integrated circuit **2400** includes one or more application processor(s) **2405** (e.g.,

CPUs), at least one graphics processor **2410**, and may additionally include an image processor **2415** and/or a video processor **2420**, any of which may be a modular IP core. In at least one embodiment, integrated circuit **2400** includes peripheral or bus logic including a USB controller **2425**, UART controller **2430**, an SPI/SDIO controller **2435**, and an I.sup.2S/I.sup.2C controller **2440**. In at least one embodiment, integrated circuit **2400** can include a display device **2445** coupled to one or more of a high-definition multimedia interface (HDMI) controller **2450** and a mobile industry processor interface (MIPI) display interface **2455**. In at least one embodiment, storage may be provided by a flash memory subsystem **2460** including flash memory and a flash memory controller. In at least one embodiment, memory interface may be provided via a memory controller **2465** for access to SDRAM or SRAM memory devices. In at least one embodiment, some integrated circuits additionally include an embedded security engine **2470**.

(386) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in integrated circuit **2400** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(387) In at least one embodiment, integrated circuit **2400** includes memory storing computer-readable executable instruction that, as a result of execution, causes one or more processors to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, integrated circuit **2400** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Integrated circuit **2400** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-23** and **25-43**.

(388) FIGS. **25A-25B** illustrate exemplary integrated circuits and associated graphics processors that may be fabricated using one or more IP cores, according to various embodiments described herein. In addition to what is illustrated, other logic and circuits may be included in at least one embodiment, including additional graphics processors/cores, peripheral interface controllers, or general-purpose processor cores.

(389) FIGS. **25A-25B** are block diagrams illustrating exemplary graphics processors for use within an SoC, according to embodiments described herein. FIG. **25A** illustrates an exemplary graphics processor **2510** of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to at least one embodiment. FIG. **25B** illustrates an additional exemplary graphics processor **2540** of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to at least one embodiment. In at least one embodiment, graphics processor **2510** of FIG. **25A** is a low power graphics processor core. In at least one embodiment, graphics processor **2540** of FIG. **25B** is a higher performance graphics processor core. In at least one embodiment, each of graphics processors **2510**, **2540** can be variants of graphics processor **2410** of FIG. **24**.

(390) In at least one embodiment, graphics processor **2510** includes a vertex processor **2505** and one or more fragment processor(s) **2515A-2515N** (e.g., **2515A**, **2515B**, **2515C**, **2515D**, through **2515N-1**, and **2515N**). In at least one embodiment, graphics processor **2510** can execute different shader programs via separate logic, such that vertex processor **2505** is optimized to execute operations for vertex shader programs, while one or more fragment processor(s) **2515A-2515N** execute fragment (e.g., pixel) shading operations for fragment or pixel shader programs. In at least one embodiment, vertex processor **2505** performs a vertex processing stage of a 3D graphics

pipeline and generates primitives and vertex data. In at least one embodiment, fragment processor(s) **2515A-2515N** use primitive and vertex data generated by vertex processor **2505** to produce a framebuffer that is displayed on a display device. In at least one embodiment, fragment processor(s) **2515A-2515N** are optimized to execute fragment shader programs as provided for in an OpenGL API, which may be used to perform similar operations as a pixel shader program as provided for in a Direct 3D API.

(391) In at least one embodiment, graphics processor **2510** additionally includes one or more memory management units (MMUs) **2520A-2520B**, cache(s) **2525A-2525B**, and circuit interconnect(s) **2530A-2530B**. In at least one embodiment, one or more MMU(s) **2520A-2520B** provide for virtual to physical address mapping for graphics processor **2510**, including for vertex processor **2505** and/or fragment processor(s) **2515A-2515N**, which may reference vertex or image/texture data stored in memory, in addition to vertex or image/texture data stored in one or more cache(s) **2525A-2525B**. In at least one embodiment, one or more MMU(s) **2520A-2520B** may be synchronized with other MMUs within system, including one or more MMUs associated with one or more application processor(s) **2405**, image processors **2415**, and/or video processors **2420** of FIG. **24**, such that each processor **2405-2420** can participate in a shared or unified virtual memory system. In at least one embodiment, one or more circuit interconnect(s) **2530A-2530B** enable graphics processor **2510** to interface with other IP cores within SoC, either via an internal bus of SoC or via a direct connection.

(392) In at least one embodiment, graphics processor **2540** includes one or more MMU(s) **2520A-2520B**, caches **2525A-2525B**, and circuit interconnects **2530A-2530B** of graphics processor **2510** of FIG. **25A**. In at least one embodiment, graphics processor **2540** includes one or more shader core(s) **2555A-2555N** (e.g., **2555A**, **2555B**, **2555C**, **2555D**, **2555E**, **2555F**, through **2555N-1**, and **2555N**), which provides for a unified shader core architecture in which a single core or type or core can execute all types of programmable shader code, including shader program code to implement vertex shaders, fragment shaders, and/or compute shaders. In at least one embodiment, a number of shader cores can vary. In at least one embodiment, graphics processor **2540** includes an inter-core task manager **2545**, which acts as a thread dispatcher to dispatch execution threads to one or more shader cores **2555A-2555N** and a tiling unit **2558** to accelerate tiling operations for tile-based rendering, in which rendering operations for a scene are subdivided in image space, for example to exploit local spatial coherence within a scene or to optimize use of internal caches.

(393) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in integrated circuit **25A** and/or **25B** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(394) In at least one embodiment, graphics processor **2540** may execute computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, graphics processor **2540** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Graphics processor **2540** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-23** and **25-43**.

(395) FIGS. **26A-26B** illustrate additional exemplary graphics processor logic according to embodiments described herein. FIG. **26A** illustrates a graphics core **2600** that may be included

within graphics processor **2410** of FIG. **24**, in at least one embodiment, and may be a unified shader core **2555A-2555N** as in FIG. **25B** in at least one embodiment. FIG. **26B** illustrates a highly-parallel general-purpose graphics processing unit **2630** suitable for deployment on a multi-chip module in at least one embodiment.

(396) In at least one embodiment, graphics core **2600** includes a shared instruction cache **2602**, a texture unit **2618**, and a cache/shared memory **2620** that are common to execution resources within graphics core **2600**. In at least one embodiment, graphics core **2600** can include multiple slices **2601A-2601N** or partition for each core, and a graphics processor can include multiple instances of graphics core **2600**. Slices **2601A-2601N** can include support logic including a local instruction cache **2604A-2604N**, a thread scheduler **2606A-2606N**, a thread dispatcher **2608A-2608N**, and a set of registers **2610A-2610N**. In at least one embodiment, slices **2601A-2601N** can include a set of additional function units (AFUs **2612A-2612N**), floating-point units (FPU **2614A-2614N**), integer arithmetic logic units (ALUs **2616-2616N**), address computational units (ACU **2613A-2613N**), double-precision floating-point units (DPFPU **2615A-2615N**), and matrix processing units (MPU **2617A-2617N**).

(397) In at least one embodiment, FPUs **2614A-2614N** can perform single-precision (32-bit) and half-precision (16-bit) floating point operations, while DPFPU **2615A-2615N** perform double precision (64-bit) floating point operations. In at least one embodiment, ALUs **2616A-2616N** can perform variable precision integer operations at 8-bit, 16-bit, and 32-bit precision, and can be configured for mixed precision operations. In at least one embodiment, MPUs **2617A-2617N** can also be configured for mixed precision matrix operations, including half-precision floating point and 8-bit integer operations. In at least one embodiment, MPUs **2617-2617N** can perform a variety of matrix operations to accelerate machine learning application frameworks, including enabling support for accelerated general matrix to matrix multiplication (GEMM). In at least one embodiment, AFUs **2612A-2612N** can perform additional logic operations not supported by floating-point or integer units, including trigonometric operations (e.g., Sine, Cosine, etc.).

(398) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in graphics core **2600** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(399) In at least one embodiment, graphics core **2600** executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, graphics core **2600** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Graphics core **2600** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-25** and **27-43**.

(400) FIG. **26B** illustrates a general-purpose processing unit (GPGPU) **2630** that can be configured to enable highly-parallel compute operations to be performed by an array of graphics processing units, in at least one embodiment. In at least one embodiment, GPGPU **2630** can be linked directly to other instances of GPGPU **2630** to create a multi-GPU cluster to improve training speed for deep neural networks. In at least one embodiment, GPGPU **2630** includes a host interface **2632** to enable a connection with a host processor. In at least one embodiment, host interface **2632** is a PCI Express interface. In at least one embodiment, host interface **2632** can be a vendor specific communications interface or communications fabric. In at least one embodiment, GPGPU **2630**

receives commands from a host processor and uses a global scheduler **2634** to distribute execution threads associated with those commands to a set of compute clusters **2636A-2636H**. In at least one embodiment, compute clusters **2636A-2636H** share a cache memory **2638**. In at least one embodiment, cache memory **2638** can serve as a higher-level cache for cache memories within compute clusters **2636A-2636H**.

(401) In at least one embodiment, GPGPU **2630** includes memory **2644A-2644B** coupled with compute clusters **2636A-2636H** via a set of memory controllers **2642A-2642B**. In at least one embodiment, memory **2644A-2644B** can include various types of memory devices including dynamic random access memory (DRAM) or graphics random access memory, such as synchronous graphics random access memory (SGRAM), including graphics double data rate (GDDR) memory.

(402) In at least one embodiment, compute clusters **2636A-2636H** each include a set of graphics cores, such as graphics core **2600** of FIG. **26A**, which can include multiple types of integer and floating point logic units that can perform computational operations at a range of precisions including suited for machine learning computations. For example, in at least one embodiment, at least a subset of floating point units in each of compute clusters **2636A-2636H** can be configured to perform 16-bit or 32-bit floating point operations, while a different subset of floating point units can be configured to perform 64-bit floating point operations.

(403) In at least one embodiment, multiple instances of GPGPU **2630** can be configured to operate as a compute cluster. In at least one embodiment, communication used by compute clusters **2636A-2636H** for synchronization and data exchange varies across embodiments. In at least one embodiment, multiple instances of GPGPU **2630** communicate over host interface **2632**. In at least one embodiment, GPGPU **2630** includes an I/O hub **2639** that couples GPGPU **2630** with a GPU link **2640** that enables a direct connection to other instances of GPGPU **2630**. In at least one embodiment, GPU link **2640** is coupled to a dedicated GPU-to-GPU bridge that enables communication and synchronization between multiple instances of GPGPU **2630**. In at least one embodiment GPU link **2640** couples with a high speed interconnect to transmit and receive data to other GPGPUs or parallel processors. In at least one embodiment, multiple instances of GPGPU **2630** are located in separate data processing systems and communicate via a network device that is accessible via host interface **2632**. In at least one embodiment GPU link **2640** can be configured to enable a connection to a host processor in addition to or as an alternative to host interface **2632**.

(404) In at least one embodiment, GPGPU **2630** can be configured to train neural networks. In at least one embodiment, GPGPU **2630** can be used within a inferencing platform. In at least one embodiment, in which GPGPU **2630** is used for inferencing, GPGPU may include fewer compute clusters **2636A-2636H** relative to when GPGPU is used for training a neural network. In at least one embodiment, memory technology associated with memory **2644A-2644B** may differ between inferencing and training configurations, with higher bandwidth memory technologies devoted to training configurations. In at least one embodiment, inferencing configuration of GPGPU **2630** can support inferencing specific instructions. For example, in at least one embodiment, an inferencing configuration can provide support for one or more 8-bit integer dot product instructions, which may be used during inferencing operations for deployed neural networks.

(405) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in GPGPU **2630** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(406) In at least one embodiment, GPGPU **2630** executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or

more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, GPGPU **2630** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Graphics core **2600** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-25** and **27-43**.

(407) FIG. **27** is a block diagram illustrating a computing system **2700** according to at least one embodiment. In at least one embodiment, computing system **2700** includes a processing subsystem **2701** having one or more processor(s) **2702** and a system memory **2704** communicating via an interconnection path that may include a memory hub **2705**. In at least one embodiment, memory hub **2705** may be a separate component within a chipset component or may be integrated within one or more processor(s) **2702**. In at least one embodiment, memory hub **2705** couples with an I/O subsystem **2711** via a communication link **2706**. In at least one embodiment, I/O subsystem **2711** includes an I/O hub **2707** that can enable computing system **2700** to receive input from one or more input device(s) **2708**. In at least one embodiment, I/O hub **2707** can enable a display controller, which may be included in one or more processor(s) **2702**, to provide outputs to one or more display device(s) **2710A**. In at least one embodiment, one or more display device(s) **2710A** coupled with I/O hub **2707** can include a local, internal, or embedded display device.

(408) In at least one embodiment, processing subsystem **2701** includes one or more parallel processor(s) **2712** coupled to memory hub **2705** via a bus or other communication link **2713**. In at least one embodiment, communication link **2713** may be one of any number of standards based communication link technologies or protocols, such as, but not limited to PCI Express, or may be a vendor specific communications interface or communications fabric. In at least one embodiment, one or more parallel processor(s) **2712** form a computationally focused parallel or vector processing system that can include a large number of processing cores and/or processing clusters, such as a many integrated core (MIC) processor. In at least one embodiment, one or more parallel processor(s) **2712** form a graphics processing subsystem that can output pixels to one of one or more display device(s) **2710A** coupled via I/O Hub **2707**. In at least one embodiment, one or more parallel processor(s) **2712** can also include a display controller and display interface (not shown) to enable a direct connection to one or more display device(s) **2710B**.

(409) In at least one embodiment, a system storage unit **2714** can connect to I/O hub **2707** to provide a storage mechanism for computing system **2700**. In at least one embodiment, an I/O switch **2716** can be used to provide an interface mechanism to enable connections between I/O hub **2707** and other components, such as a network adapter **2718** and/or wireless network adapter **2719** that may be integrated into platform, and various other devices that can be added via one or more add-in device(s) **2720**. In at least one embodiment, network adapter **2718** can be an Ethernet adapter or another wired network adapter. In at least one embodiment, wireless network adapter **2719** can include one or more of a Wi-Fi, Bluetooth, near field communication (NFC), or other network device that includes one or more wireless radios.

(410) In at least one embodiment, computing system **2700** can include other components not explicitly shown, including USB or other port connections, optical storage drives, video capture devices, and like, may also be connected to I/O hub **2707**. In at least one embodiment, communication paths interconnecting various components in FIG. **27** may be implemented using any suitable protocols, such as PCI (Peripheral Component Interconnect) based protocols (e.g., PCI-Express), or other bus or point-to-point communication interfaces and/or protocol(s), such as NV-Link high-speed interconnect, or interconnect protocols.

(411) In at least one embodiment, one or more parallel processor(s) **2712** incorporate circuitry optimized for graphics and video processing, including, for example, video output circuitry, and constitutes a graphics processing unit (GPU). In at least one embodiment, one or more parallel

processor(s) **2712** incorporate circuitry optimized for general purpose processing. In at least one embodiment, components of computing system **2700** may be integrated with one or more other system elements on a single integrated circuit. For example, in at least one embodiment, one or more parallel processor(s) **2712**, memory hub **2705**, processor(s) **2702**, and I/O hub **2707** can be integrated into a system on chip (SoC) integrated circuit. In at least one embodiment, components of computing system **2700** can be integrated into a single package to form a system in package (SIP) configuration. In at least one embodiment, at least a portion of components of computing system **2700** can be integrated into a multi-chip module (MCM), which can be interconnected with other multi-chip modules into a modular computing system.

(412) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in system FIG. **2700** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(413) In at least one embodiment, computing system **2700** or a component thereof (e.g., one or more parallel processor(s) **2712**) executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, computing system **2700** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Computing system **2700** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-26** and **28-43**.

(414) Processors

(415) FIG. **28A** illustrates a parallel processor **2800** according to at least one embodiment. In at least one embodiment, various components of parallel processor **2800** may be implemented using one or more integrated circuit devices, such as programmable processors, application specific integrated circuits (ASICs), or field programmable gate arrays (FPGA). In at least one embodiment, illustrated parallel processor **2800** is a variant of one or more parallel processor(s) **2712** shown in FIG. **27** according to an exemplary embodiment.

(416) In at least one embodiment, parallel processor **2800** includes a parallel processing unit **2802**. In at least one embodiment, parallel processing unit **2802** includes an I/O unit **2804** that enables communication with other devices, including other instances of parallel processing unit **2802**. In at least one embodiment, I/O unit **2804** may be directly connected to other devices. In at least one embodiment, I/O unit **2804** connects with other devices via use of a hub or switch interface, such as memory hub **2705**. In at least one embodiment, connections between memory hub **2705** and I/O unit **2804** form a communication link **2713**. In at least one embodiment, I/O unit **2804** connects with a host interface **2806** and a memory crossbar **2816**, where host interface **2806** receives commands directed to performing processing operations and memory crossbar **2816** receives commands directed to performing memory operations.

(417) In at least one embodiment, when host interface **2806** receives a command buffer via I/O unit **2804**, host interface **2806** can direct work operations to perform those commands to a front end **2808**. In at least one embodiment, front end **2808** couples with a scheduler **2810**, which is configured to distribute commands or other work items to a processing cluster array **2812**. In at least one embodiment, scheduler **2810** ensures that processing cluster array **2812** is properly configured and in a valid state before tasks are distributed to processing cluster array **2812** of processing cluster array **2812**. In at least one embodiment, scheduler **2810** is implemented via

firmware logic executing on a microcontroller. In at least one embodiment, microcontroller implemented scheduler **2810** is configurable to perform complex scheduling and work distribution operations at coarse and fine granularity, enabling rapid preemption and context switching of threads executing on processing array **2812**. In at least one embodiment, host software can provide workloads for scheduling on processing array **2812** via one of multiple graphics processing doorbells. In at least one embodiment, workloads can then be automatically distributed across processing array **2812** by scheduler **2810** logic within a microcontroller including scheduler **2810**. (418) In at least one embodiment, processing cluster array **2812** can include up to “N” processing clusters (e.g., cluster **2814A**, cluster **2814B**, through cluster **2814N**). In at least one embodiment, each cluster **2814A-2814N** of processing cluster array **2812** can execute a large number of concurrent threads. In at least one embodiment, scheduler **2810** can allocate work to clusters **2814A-2814N** of processing cluster array **2812** using various scheduling and/or work distribution algorithms, which may vary depending on workload arising for each type of program or computation. In at least one embodiment, scheduling can be handled dynamically by scheduler **2810**, or can be assisted in part by compiler logic during compilation of program logic configured for execution by processing cluster array **2812**. In at least one embodiment, different clusters **2814A-2814N** of processing cluster array **2812** can be allocated for processing different types of programs or for performing different types of computations.

(419) In at least one embodiment, processing cluster array **2812** can be configured to perform various types of parallel processing operations. In at least one embodiment, processing cluster array **2812** is configured to perform general-purpose parallel compute operations. For example, in at least one embodiment, processing cluster array **2812** can include logic to execute processing tasks including filtering of video and/or audio data, performing modeling operations, including physics operations, and performing data transformations.

(420) In at least one embodiment, processing cluster array **2812** is configured to perform parallel graphics processing operations. In at least one embodiment, processing cluster array **2812** can include additional logic to support execution of such graphics processing operations, including, but not limited to texture sampling logic to perform texture operations, as well as tessellation logic and other vertex processing logic. In at least one embodiment, processing cluster array **2812** can be configured to execute graphics processing related shader programs such as, but not limited to vertex shaders, tessellation shaders, geometry shaders, and pixel shaders. In at least one embodiment, parallel processing unit **2802** can transfer data from system memory via I/O unit **2804** for processing. In at least one embodiment, during processing, transferred data can be stored to on-chip memory (e.g., parallel processor memory **2822**) during processing, then written back to system memory.

(421) In at least one embodiment, when parallel processing unit **2802** is used to perform graphics processing, scheduler **2810** can be configured to divide a processing workload into approximately equal sized tasks, to better enable distribution of graphics processing operations to multiple clusters **2814A-2814N** of processing cluster array **2812**. In at least one embodiment, portions of processing cluster array **2812** can be configured to perform different types of processing. For example, in at least one embodiment, a first portion may be configured to perform vertex shading and topology generation, a second portion may be configured to perform tessellation and geometry shading, and a third portion may be configured to perform pixel shading or other screen space operations, to produce a rendered image for display. In at least one embodiment, intermediate data produced by one or more of clusters **2814A-2814N** may be stored in buffers to allow intermediate data to be transmitted between clusters **2814A-2814N** for further processing.

(422) In at least one embodiment, processing cluster array **2812** can receive processing tasks to be executed via scheduler **2810**, which receives commands defining processing tasks from front end **2808**. In at least one embodiment, processing tasks can include indices of data to be processed, e.g., surface (patch) data, primitive data, vertex data, and/or pixel data, as well as state parameters and

commands defining how data is to be processed (e.g., what program is to be executed). In at least one embodiment, scheduler **2810** may be configured to fetch indices corresponding to tasks or may receive indices from front end **2808**. In at least one embodiment, front end **2808** can be configured to ensure processing cluster array **2812** is configured to a valid state before a workload specified by incoming command buffers (e.g., batch-buffers, push buffers, etc.) is initiated.

(423) In at least one embodiment, each of one or more instances of parallel processing unit **2802** can couple with parallel processor memory **2822**. In at least one embodiment, parallel processor memory **2822** can be accessed via memory crossbar **2816**, which can receive memory requests from processing cluster array **2812** as well as I/O unit **2804**. In at least one embodiment, memory crossbar **2816** can access parallel processor memory **2822** via a memory interface **2818**. In at least one embodiment, memory interface **2818** can include multiple partition units (e.g., partition unit **2820A**, partition unit **2820B**, through partition unit **2820N**) that can each couple to a portion (e.g., memory unit) of parallel processor memory **2822**. In at least one embodiment, a number of partition units **2820A-2820N** is configured to be equal to a number of memory units, such that a first partition unit **2820A** has a corresponding first memory unit **2824A**, a second partition unit **2820B** has a corresponding memory unit **2824B**, and an Nth partition unit **2820N** has a corresponding Nth memory unit **2824N**. In at least one embodiment, a number of partition units **2820A-2820N** may not be equal to a number of memory devices.

(424) In at least one embodiment, memory units **2824A-2824N** can include various types of memory devices, including dynamic random access memory (DRAM) or graphics random access memory, such as synchronous graphics random access memory (SGRAM), including graphics double data rate (GDDR) memory. In at least one embodiment, memory units **2824A-2824N** may also include 3D stacked memory, including but not limited to high bandwidth memory (HBM). In at least one embodiment, render targets, such as frame buffers or texture maps may be stored across memory units **2824A-2824N**, allowing partition units **2820A-2820N** to write portions of each render target in parallel to efficiently use available bandwidth of parallel processor memory **2822**. In at least one embodiment, a local instance of parallel processor memory **2822** may be excluded in favor of a unified memory design that utilizes system memory in conjunction with local cache memory.

(425) In at least one embodiment, any one of clusters **2814A-2814N** of processing cluster array **2812** can process data that will be written to any of memory units **2824A-2824N** within parallel processor memory **2822**. In at least one embodiment, memory crossbar **2816** can be configured to transfer an output of each cluster **2814A-2814N** to any partition unit **2820A-2820N** or to another cluster **2814A-2814N**, which can perform additional processing operations on an output. In at least one embodiment, each cluster **2814A-2814N** can communicate with memory interface **2818** through memory crossbar **2816** to read from or write to various external memory devices. In at least one embodiment, memory crossbar **2816** has a connection to memory interface **2818** to communicate with I/O unit **2804**, as well as a connection to a local instance of parallel processor memory **2822**, enabling processing units within different processing clusters **2814A-2814N** to communicate with system memory or other memory that is not local to parallel processing unit **2802**. In at least one embodiment, memory crossbar **2816** can use virtual channels to separate traffic streams between clusters **2814A-2814N** and partition units **2820A-2820N**.

(426) In at least one embodiment, multiple instances of parallel processing unit **2802** can be provided on a single add-in card, or multiple add-in cards can be interconnected. In at least one embodiment, different instances of parallel processing unit **2802** can be configured to inter-operate even if different instances have different numbers of processing cores, different amounts of local parallel processor memory, and/or other configuration differences. For example, in at least one embodiment, some instances of parallel processing unit **2802** can include higher precision floating point units relative to other instances. In at least one embodiment, systems incorporating one or more instances of parallel processing unit **2802** or parallel processor **2800** can be implemented in a

variety of configurations and form factors, including but not limited to desktop, laptop, or handheld personal computers, servers, workstations, game consoles, and/or embedded systems.

(427) FIG. **28B** is a block diagram of a partition unit **2820** according to at least one embodiment. In at least one embodiment, partition unit **2820** is an instance of one of partition units **2820A-2820N** of FIG. **28A**. In at least one embodiment, partition unit **2820** includes an L2 cache **2821**, a frame buffer interface **2825**, and a ROP **2826** (raster operations unit). L2 cache **2821** is a read/write cache that is configured to perform load and store operations received from memory crossbar **2816** and ROP **2826**. In at least one embodiment, read misses and urgent write-back requests are output by L2 cache **2821** to frame buffer interface **2825** for processing. In at least one embodiment, updates can also be sent to a frame buffer via frame buffer interface **2825** for processing. In at least one embodiment, frame buffer interface **2825** interfaces with one of memory units in parallel processor memory, such as memory units **2824A-2824N** of FIG. **28** (e.g., within parallel processor memory **2822**).

(428) In at least one embodiment, ROP **2826** is a processing unit that performs raster operations such as stencil, z test, blending, and like. In at least one embodiment, ROP **2826** then outputs processed graphics data that is stored in graphics memory. In at least one embodiment, ROP **2826** includes compression logic to compress depth or color data that is written to memory and decompress depth or color data that is read from memory. In at least one embodiment, compression logic can be lossless compression logic that makes use of one or more of multiple compression algorithms. type of compression that is performed by ROP **2826** can vary based on statistical characteristics of data to be compressed. For example, in at least one embodiment, delta color compression is performed on depth and color data on a per-tile basis.

(429) In In at least one embodiment, ROP **2826** is included within each processing cluster (e.g., cluster **2814A-2814N** of FIG. **28**) instead of within partition unit **2820**. In at least one embodiment, read and write requests for pixel data are transmitted over memory crossbar **2816** instead of pixel fragment data. In at least one embodiment, processed graphics data may be displayed on a display device, such as one of one or more display device(s) **2710** of FIG. **27**, routed for further processing by processor(s) **2702**, or routed for further processing by one of processing entities within parallel processor **2800** of FIG. **28A**.

(430) FIG. **28C** is a block diagram of a processing cluster **2814** within a parallel processing unit according to at least one embodiment. In at least one embodiment, a processing cluster is an instance of one of processing clusters **2814A-2814N** of FIG. **28**. In at least one embodiment, processing cluster **2814** can be configured to execute many threads in parallel, where term “thread” refers to an instance of a particular program executing on a particular set of input data. In at least one embodiment, single-instruction, multiple-data (SIMD) instruction issue techniques are used to support parallel execution of a large number of threads without providing multiple independent instruction units. In at least one embodiment, single-instruction, multiple-thread (SIMT) techniques are used to support parallel execution of a large number of generally synchronized threads, using a common instruction unit configured to issue instructions to a set of processing engines within each one of processing clusters.

(431) In at least one embodiment, operation of processing cluster **2814** can be controlled via a pipeline manager **2832** that distributes processing tasks to SIMT parallel processors. In at least one embodiment, pipeline manager **2832** receives instructions from scheduler **2810** of FIG. **28** and manages execution of those instructions via a graphics multiprocessor **2834** and/or a texture unit **2836**. In at least one embodiment, graphics multiprocessor **2834** is an exemplary instance of a SIMT parallel processor. However, in at least one embodiment, various types of SIMT parallel processors of differing architectures may be included within processing cluster **2814**. In at least one embodiment, one or more instances of graphics multiprocessor **2834** can be included within a processing cluster **2814**. In at least one embodiment, graphics multiprocessor **2834** can process data and a data crossbar **2840** can be used to distribute processed data to one of multiple possible

destinations, including other shader units. In at least one embodiment, pipeline manager **2832** can facilitate distribution of processed data by specifying destinations for processed data to be distributed via data crossbar **2840**.

(432) In at least one embodiment, each graphics multiprocessor **2834** within processing cluster **2814** can include an identical set of functional execution logic (e.g., arithmetic logic units, load-store units, etc.). In at least one embodiment, functional execution logic can be configured in a pipelined manner in which new instructions can be issued before previous instructions are complete. In at least one embodiment, functional execution logic supports a variety of operations including integer and floating point arithmetic, comparison operations, Boolean operations, bit-shifting, and computation of various algebraic functions. In at least one embodiment, same functional-unit hardware can be leveraged to perform different operations and any combination of functional units may be present.

(433) In at least one embodiment, instructions transmitted to processing cluster **2814** constitute a thread. In at least one embodiment, a set of threads executing across a set of parallel processing engines is a thread group. In at least one embodiment, thread group executes a program on different input data. In at least one embodiment, each thread within a thread group can be assigned to a different processing engine within a graphics multiprocessor **2834**. In at least one embodiment, a thread group may include fewer threads than a number of processing engines within graphics multiprocessor **2834**. In at least one embodiment, when a thread group includes fewer threads than a number of processing engines, one or more of processing engines may be idle during cycles in which that thread group is being processed. In at least one embodiment, a thread group may also include more threads than a number of processing engines within graphics multiprocessor **2834**. In at least one embodiment, when a thread group includes more threads than number of processing engines within graphics multiprocessor **2834**, processing can be performed over consecutive clock cycles. In at least one embodiment, multiple thread groups can be executed concurrently on a graphics multiprocessor **2834**.

(434) In at least one embodiment, graphics multiprocessor **2834** includes an internal cache memory to perform load and store operations. In at least one embodiment, graphics multiprocessor **2834** can forego an internal cache and use a cache memory (e.g., L1 cache **2848**) within processing cluster **2814**. In at least one embodiment, each graphics multiprocessor **2834** also has access to L2 caches within partition units (e.g., partition units **2820A-2820N** of FIG. **28**) that are shared among all processing clusters **2814** and may be used to transfer data between threads. In at least one embodiment, graphics multiprocessor **2834** may also access off-chip global memory, which can include one or more of local parallel processor memory and/or system memory. In at least one embodiment, any memory external to parallel processing unit **2802** may be used as global memory. In at least one embodiment, processing cluster **2814** includes multiple instances of graphics multiprocessor **2834** can share common instructions and data, which may be stored in L1 cache **2848**.

(435) In at least one embodiment, each processing cluster **2814** may include an MMU **2845** (memory management unit) that is configured to map virtual addresses into physical addresses. In at least one embodiment, one or more instances of MMU **2845** may reside within memory interface **2818** of FIG. **28**. In at least one embodiment, MMU **2845** includes a set of page table entries (PTEs) used to map a virtual address to a physical address of a tile (talk more about tiling) and optionally a cache line index. In at least one embodiment, MMU **2845** may include address translation lookaside buffers (TLB) or caches that may reside within graphics multiprocessor **2834** or L1 cache or processing cluster **2814**. In at least one embodiment, physical address is processed to distribute surface data access locality to allow efficient request interleaving among partition units. In at least one embodiment, cache line index may be used to determine whether a request for a cache line is a hit or miss.

(436) In at least one embodiment, a processing cluster **2814** may be configured such that each

graphics multiprocessor **2834** is coupled to a texture unit **2836** for performing texture mapping operations, e.g., determining texture sample positions, reading texture data, and filtering texture data. In at least one embodiment, texture data is read from an internal texture L1 cache (not shown) or from an L1 cache within graphics multiprocessor **2834** and is fetched from an L2 cache, local parallel processor memory, or system memory, as needed. In at least one embodiment, each graphics multiprocessor **2834** outputs processed tasks to data crossbar **2840** to provide processed task to another processing cluster **2814** for further processing or to store processed task in an L2 cache, local parallel processor memory, or system memory via memory crossbar **2816**. In at least one embodiment, preROP **2842** (pre-raster operations unit) is configured to receive data from graphics multiprocessor **2834**, direct data to ROP units, which may be located with partition units as described herein (e.g., partition units **2820A-2820N** of FIG. **28**). In at least one embodiment, PreROP **2842** unit can perform optimizations for color blending, organize pixel color data, and perform address translations.

(437) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in graphics processing cluster **2814** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(438) In at least one embodiment, parallel processor **2800** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, parallel processor **2800** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Parallel processor **2800** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-27** and **29-43**.

(439) FIG. **28D** shows a graphics multiprocessor **2834** according to at least one embodiment. In at least one embodiment, graphics multiprocessor **2834** couples with pipeline manager **2832** of processing cluster **2814**. In at least one embodiment, graphics multiprocessor **2834** has an execution pipeline including but not limited to an instruction cache **2852**, an instruction unit **2854**, an address mapping unit **2856**, a register file **2858**, one or more general purpose graphics processing unit (GPGPU) cores **2862**, and one or more load/store units **2866**. GPGPU cores **2862** and load/store units **2866** are coupled with cache memory **2872** and shared memory **2870** via a memory and cache interconnect **2868**.

(440) In at least one embodiment, instruction cache **2852** receives a stream of instructions to execute from pipeline manager **2832**. In at least one embodiment, instructions are cached in instruction cache **2852** and dispatched for execution by instruction unit **2854**. In at least one embodiment, instruction unit **2854** can dispatch instructions as thread groups (e.g., warps), with each thread of thread group assigned to a different execution unit within GPGPU core **2862**. In at least one embodiment, an instruction can access any of a local, shared, or global address space by specifying an address within a unified address space. In at least one embodiment, address mapping unit **2856** can be used to translate addresses in a unified address space into a distinct memory address that can be accessed by load/store units **2866**.

(441) In at least one embodiment, register file **2858** provides a set of registers for functional units of graphics multiprocessor **2834**. In at least one embodiment, register file **2858** provides temporary storage for operands connected to data paths of functional units (e.g., GPGPU cores **2862**,

load/store units **2866**) of graphics multiprocessor **2834**. In at least one embodiment, register file **2858** is divided between each of functional units such that each functional unit is allocated a dedicated portion of register file **2858**. In at least one embodiment, register file **2858** is divided between different warps being executed by graphics multiprocessor **2834**.

(442) In at least one embodiment, GPGPU cores **2862** can each include floating point units (FPUs) and/or integer arithmetic logic units (ALUs) that are used to execute instructions of graphics multiprocessor **2834**. GPGPU cores **2862** can be similar in architecture or can differ in architecture. In at least one embodiment, a first portion of GPGPU cores **2862** include a single precision FPU and an integer ALU while a second portion of GPGPU cores include a double precision FPU. In at least one embodiment, FPUs can implement IEEE 754-2008 standard for floating point arithmetic or enable variable precision floating point arithmetic. In at least one embodiment, graphics multiprocessor **2834** can additionally include one or more fixed function or special function units to perform specific functions such as copy rectangle or pixel blending operations. In at least one embodiment one or more of GPGPU cores can also include fixed or special function logic.

(443) In at least one embodiment, GPGPU cores **2862** include SIMD logic capable of performing a single instruction on multiple sets of data. In at least one embodiment GPGPU cores **2862** can physically execute SIMD4, SIMD8, and SIMD16 instructions and logically execute SIMD1, SIMD2, and SIMD32 instructions. In at least one embodiment, SIMD instructions for GPGPU cores can be generated at compile time by a shader compiler or automatically generated when executing programs written and compiled for single program multiple data (SPMD) or SIMT architectures. In at least one embodiment, multiple threads of a program configured for an SIMT execution model can be executed via a single SIMD instruction. For example, in at least one embodiment, eight SIMT threads that perform same or similar operations can be executed in parallel via a single SIMD8 logic unit.

(444) In at least one embodiment, memory and cache interconnect **2868** is an interconnect network that connects each functional unit of graphics multiprocessor **2834** to register file **2858** and to shared memory **2870**. In at least one embodiment, memory and cache interconnect **2868** is a crossbar interconnect that allows load/store unit **2866** to implement load and store operations between shared memory **2870** and register file **2858**. In at least one embodiment, register file **2858** can operate at a same frequency as GPGPU cores **2862**, thus data transfer between GPGPU cores **2862** and register file **2858** is very low latency. In at least one embodiment, shared memory **2870** can be used to enable communication between threads that execute on functional units within graphics multiprocessor **2834**. In at least one embodiment, cache memory **2872** can be used as a data cache for example, to cache texture data communicated between functional units and texture unit **2836**. In at least one embodiment, shared memory **2870** can also be used as a program managed cache. In at least one embodiment, threads executing on GPGPU cores **2862** can programmatically store data within shared memory in addition to automatically cached data that is stored within cache memory **2872**.

(445) In at least one embodiment, a parallel processor or GPGPU as described herein is communicatively coupled to host/processor cores to accelerate graphics operations, machine-learning operations, pattern analysis operations, and various general purpose GPU (GPGPU) functions. In at least one embodiment, GPU may be communicatively coupled to host processor/cores over a bus or other interconnect (e.g., a high speed interconnect such as PCIe or NVLink). In at least one embodiment, GPU may be integrated on same package or chip as cores and communicatively coupled to cores over an internal processor bus/interconnect (i.e., internal to package or chip). In at least one embodiment, regardless of manner in which GPU is connected, processor cores may allocate work to GPU in form of sequences of commands/instructions contained in a work descriptor. In at least one embodiment, GPU then uses dedicated circuitry/logic for efficiently processing these commands/instructions.

(446) Inference and/or training logic **1515** are used to perform inferencing and/or training

operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in graphics multiprocessor **2834** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(447) In at least one embodiment, graphics multiprocessor **2834** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, graphics multiprocessor **2834** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Graphics multiprocessor **2834** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-27** and **29-43**.

(448) FIG. **29** illustrates a multi-GPU computing system **2900**, according to at least one embodiment. In at least one embodiment, multi-GPU computing system **2900** can include a processor **2902** coupled to multiple general purpose graphics processing units (GPGPUs) **2906A-D** via a host interface switch **2904**. In at least one embodiment, host interface switch **2904** is a PCI express switch device that couples processor **2902** to a PCI express bus over which processor **2902** can communicate with GPGPUs **2906A-D**. GPGPUs **2906A-D** can interconnect via a set of high-speed point to point GPU to GPU links **2916**. In at least one embodiment, GPU to GPU links **2916** connect to each of GPGPUs **2906A-D** via a dedicated GPU link. In at least one embodiment, P2P GPU links **2916** enable direct communication between each of GPGPUs **2906A-D** without requiring communication over host interface bus **2904** to which processor **2902** is connected. In at least one embodiment, with GPU-to-GPU traffic directed to P2P GPU links **2916**, host interface bus **2904** remains available for system memory access or to communicate with other instances of multi-GPU computing system **2900**, for example, via one or more network devices. While in at least one embodiment GPGPUs **2906A-D** connect to processor **2902** via host interface switch **2904**, in at least one embodiment processor **2902** includes direct support for P2P GPU links **2916** and can connect directly to GPGPUs **2906A-D**.

(449) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in multi-GPU computing system **2900** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(450) In at least one embodiment, multi-GPU computing system **2900** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, multi-GPU computing system **2900** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Multi-GPU computing system **2900** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-28** and **30-43**.

(451) FIG. **30** is a block diagram of a graphics processor **3000**, according to at least one

embodiment. In at least one embodiment, graphics processor **3000** includes a ring interconnect **3002**, a pipeline front-end **3004**, a media engine **3037**, and graphics cores **3080A-3080N**. In at least one embodiment, ring interconnect **3002** couples graphics processor **3000** to other processing units, including other graphics processors or one or more general-purpose processor cores. In at least one embodiment, graphics processor **3000** is one of many processors integrated within a multi-core processing system.

(452) In at least one embodiment, graphics processor **3000** receives batches of commands via ring interconnect **3002**. In at least one embodiment, incoming commands are interpreted by a command streamer **3003** in pipeline front-end **3004**. In at least one embodiment, graphics processor **3000** includes scalable execution logic to perform 3D geometry processing and media processing via graphics core(s) **3080A-3080N**. In at least one embodiment, for 3D geometry processing commands, command streamer **3003** supplies commands to geometry pipeline **3036**. In at least one embodiment, for at least some media processing commands, command streamer **3003** supplies commands to a video front end **3034**, which couples with a media engine **3037**. In at least one embodiment, media engine **3037** includes a Video Quality Engine (VQE) **3030** for video and image post-processing and a multi-format encode/decode (MFX) **3033** engine to provide hardware-accelerated media data encode and decode. In at least one embodiment, geometry pipeline **3036** and media engine **3037** each generate execution threads for thread execution resources provided by at least one graphics core **3080A**.

(453) In at least one embodiment, graphics processor **3000** includes scalable thread execution resources featuring modular cores **3080A-3080N** (sometimes referred to as core slices), each having multiple sub-cores **3050A-550N**, **3060A-3060N** (sometimes referred to as core sub-slices). In at least one embodiment, graphics processor **3000** can have any number of graphics cores **3080A** through **3080N**. In at least one embodiment, graphics processor **3000** includes a graphics core **3080A** having at least a first sub-core **3050A** and a second sub-core **3060A**. In at least one embodiment, graphics processor **3000** is a low power processor with a single sub-core (e.g., **3050A**). In at least one embodiment, graphics processor **3000** includes multiple graphics cores **3080A-3080N**, each including a set of first sub-cores **3050A-3050N** and a set of second sub-cores **3060A-3060N**. In at least one embodiment, each sub-core in first sub-cores **3050A-3050N** includes at least a first set of execution units **3052A-3052N** and media/texture samplers **3054A-3054N**. In at least one embodiment, each sub-core in second sub-cores **3060A-3060N** includes at least a second set of execution units **3062A-3062N** and samplers **3064A-3064N**. In at least one embodiment, each sub-core **3050A-3050N**, **3060A-3060N** shares a set of shared resources **3070A-3070N**. In at least one embodiment, shared resources include shared cache memory and pixel operation logic.

(454) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, inference and/or training logic **1515** may be used in graphics processor **3000** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(455) In at least one embodiment, graphics processor **3000** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, graphics processor **3000** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Graphics processor **3000** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-29** and

31-43.

(456) FIG. **31** is a block diagram illustrating micro-architecture for a processor **3100** that may include logic circuits to perform instructions, according to at least one embodiment. In at least one embodiment, processor **3100** may perform instructions, including x86 instructions, ARM instructions, specialized instructions for application-specific integrated circuits (ASICs), etc. In at least one embodiment, processor **3110** may include registers to store packed data, such as 64-bit wide MMX™ registers in microprocessors enabled with MMX technology from Intel Corporation of Santa Clara, Calif. In at least one embodiment, MMX registers, available in both integer and floating point forms, may operate with packed data elements that accompany single instruction, multiple data (“SIMD”) and streaming SIMD extensions (“SSE”) instructions. In at least one embodiment, 128-bit wide XMM registers relating to SSE2, SSE3, SSE4, AVX, or beyond (referred to generically as “SSEx”) technology may hold such packed data operands. In at least one embodiment, processors **3110** may perform instructions to accelerate machine learning or deep learning algorithms, training, or inferencing.

(457) In at least one embodiment, processor **3100** includes an in-order front end (“front end”) **3101** to fetch instructions to be executed and prepare instructions to be used later in processor pipeline. In at least one embodiment, front end **3101** may include several units. In at least one embodiment, an instruction prefetcher **3126** fetches instructions from memory and feeds instructions to an instruction decoder **3128** which in turn decodes or interprets instructions. For example, in at least one embodiment, instruction decoder **3128** decodes a received instruction into one or more operations called “micro-instructions” or “micro-operations” (also called “micro ops” or “uops”) that machine may execute. In at least one embodiment, instruction decoder **3128** parses instruction into an opcode and corresponding data and control fields that may be used by micro-architecture to perform operations in accordance with at least one embodiment. In at least one embodiment, a trace cache **3130** may assemble decoded uops into program ordered sequences or traces in a uop queue **3134** for execution. In at least one embodiment, when trace cache **3130** encounters a complex instruction, a microcode ROM **3132** provides uops needed to complete operation.

(458) In at least one embodiment, some instructions may be converted into a single micro-op, whereas others need several micro-ops to complete full operation. In at least one embodiment, if more than four micro-ops are needed to complete an instruction, instruction decoder **3128** may access microcode ROM **3132** to perform instruction. In at least one embodiment, an instruction may be decoded into a small number of micro-ops for processing at instruction decoder **3128**. In at least one embodiment, an instruction may be stored within microcode ROM **3132** should a number of micro-ops be needed to accomplish operation. In at least one embodiment, trace cache **3130** refers to an entry point programmable logic array (“PLA”) to determine a correct micro-instruction pointer for reading microcode sequences to complete one or more instructions from microcode ROM **3132** in accordance with at least one embodiment. In at least one embodiment, after microcode ROM **3132** finishes sequencing micro-ops for an instruction, front end **3101** of machine may resume fetching micro-ops from trace cache **3130**.

(459) In at least one embodiment, out-of-order execution engine (“out of order engine”) **3103** may prepare instructions for execution. In at least one embodiment, out-of-order execution logic has a number of buffers to smooth out and re-order flow of instructions to optimize performance as they go down pipeline and get scheduled for execution. out-of-order execution engine **3103** includes, without limitation, an allocator/register renamer **3140**, a memory uop queue **3142**, an integer/floating point uop queue **3144**, a memory scheduler **3146**, a fast scheduler **3102**, a slow/general floating point scheduler (“slow/general FP scheduler”) **3104**, and a simple floating point scheduler (“simple FP scheduler”) **3106**. In at least one embodiment, fast scheduler **3102**, slow/general floating point scheduler **3104**, and simple floating point scheduler **3106** are also collectively referred to herein as “uop schedulers **3102**, **3104**, **3106**.” allocator/register renamer **3140** allocates machine buffers and resources that each uop needs in order to execute. In at least

one embodiment, allocator/register renamer **3140** renames logic registers onto entries in a register file. In at least one embodiment, allocator/register renamer **3140** also allocates an entry for each uop in one of two uop queues, memory uop queue **3142** for memory operations and integer/floating point uop queue **3144** for non-memory operations, in front of memory scheduler **3146** and uop schedulers **3102**, **3104**, **3106**. In at least one embodiment, uop schedulers **3102**, **3104**, **3106**, determine when a uop is ready to execute based on readiness of their dependent input register operand sources and availability of execution resources uops need to complete their operation. In at least one embodiment, fast scheduler **3102** of at least one embodiment may schedule on each half of main clock cycle while slow/general floating point scheduler **3104** and simple floating point scheduler **3106** may schedule once per main processor clock cycle. In at least one embodiment, uop schedulers **3102**, **3104**, **3106** arbitrate for dispatch ports to schedule uops for execution.

(460) In at least one embodiment, execution block **b11** includes, without limitation, an integer register file/bypass network **3108**, a floating point register file/bypass network (“FP register file/bypass network”) **3110**, address generation units (“AGUs”) **3112** and **3114**, fast Arithmetic Logic Units (ALUs) (“fast ALUs”) **3116** and **3118**, a slow Arithmetic Logic Unit (“slow ALU”) **3120**, a floating point ALU (“FP”) **3122**, and a floating point move unit (“FP move”) **3124**. In at least one embodiment, integer register file/bypass network **3108** and floating point register file/bypass network **3110** are also referred to herein as “register files **3108**, **3110**.” In at least one embodiment, AGUs **3112** and **3114**, fast ALUs **3116** and **3118**, slow ALU **3120**, floating point ALU **3122**, and floating point move unit **3124** are also referred to herein as “execution units **3112**, **3114**, **3116**, **3118**, **3120**, **3122**, and **3124**.” In at least one embodiment, execution block **b11** may include, without limitation, any number (including zero) and type of register files, bypass networks, address generation units, and execution units, in any combination.

(461) In at least one embodiment, register files **3108**, **3110** may be arranged between uop schedulers **3102**, **3104**, **3106**, and execution units **3112**, **3114**, **3116**, **3118**, **3120**, **3122**, and **3124**. In at least one embodiment, integer register file/bypass network **3108** performs integer operations. In at least one embodiment, floating point register file/bypass network **3110** performs floating point operations. In at least one embodiment, each of register files **3108**, **3110** may include, without limitation, a bypass network that may bypass or forward just completed results that have not yet been written into register file to new dependent uops. In at least one embodiment, register files **3108**, **3110** may communicate data with each other. In at least one embodiment, integer register file/bypass network **3108** may include, without limitation, two separate register files, one register file for low-order thirty-two bits of data and a second register file for high order thirty-two bits of data. In at least one embodiment, floating point register file/bypass network **3110** may include, without limitation, 128-bit wide entries because floating point instructions typically have operands from 64 to 128 bits in width.

(462) In at least one embodiment, execution units **3112**, **3114**, **3116**, **3118**, **3120**, **3122**, **3124** may execute instructions. In at least one embodiment, register files **3108**, **3110** store integer and floating point data operand values that micro-instructions need to execute. In at least one embodiment, processor **3100** may include, without limitation, any number and combination of execution units **3112**, **3114**, **3116**, **3118**, **3120**, **3122**, **3124**. In at least one embodiment, floating point ALU **3122** and floating point move unit **3124**, may execute floating point, MMX, SIMD, AVX and SSE, or other operations, including specialized machine learning instructions. In at least one embodiment, floating point ALU **3122** may include, without limitation, a 64-bit by 64-bit floating point divider to execute divide, square root, and remainder micro ops. In at least one embodiment, instructions involving a floating point value may be handled with floating point hardware. In at least one embodiment, ALU operations may be passed to fast ALUs **3116**, **3118**. In at least one embodiment, fast ALUS **3116**, **3118** may execute fast operations with an effective latency of half a clock cycle. In at least one embodiment, most complex integer operations go to slow ALU **3120** as slow ALU **3120** may include, without limitation, integer execution hardware for long-latency type of

operations, such as a multiplier, shifts, flag logic, and branch processing. In at least one embodiment, memory load/store operations may be executed by AGUS **3112**, **3114**. In at least one embodiment, fast ALU **3116**, fast ALU **3118**, and slow ALU **3120** may perform integer operations on 64-bit data operands. In at least one embodiment, fast ALU **3116**, fast ALU **3118**, and slow ALU **3120** may be implemented to support a variety of data bit sizes including sixteen, thirty-two, 128, 256, etc. In at least one embodiment, floating point ALU **3122** and floating point move unit **3124** may be implemented to support a range of operands having bits of various widths. In at least one embodiment, floating point ALU **3122** and floating point move unit **3124** may operate on 128-bit wide packed data operands in conjunction with SIMD and multimedia instructions.

(463) In at least one embodiment, uop schedulers **3102**, **3104**, **3106**, dispatch dependent operations before parent load has finished executing. In at least one embodiment, as uops may be speculatively scheduled and executed in processor **3100**, processor **3100** may also include logic to handle memory misses. In at least one embodiment, if a data load misses in data cache, there may be dependent operations in flight in pipeline that have left scheduler with temporarily incorrect data. In at least one embodiment, a replay mechanism tracks and re-executes instructions that use incorrect data. In at least one embodiment, dependent operations might need to be replayed and independent ones may be allowed to complete. In at least one embodiment, schedulers and replay mechanism of at least one embodiment of a processor may also be designed to catch instruction sequences for text string comparison operations.

(464) In at least one embodiment, term “registers” may refer to on-board processor storage locations that may be used as part of instructions to identify operands. In at least one embodiment—registers may be those that may be usable from outside of processor (from a programmer's perspective). In at least one embodiment, registers might not be limited to a particular type of circuit. Rather, in at least one embodiment, a register may store data, provide data, and perform functions described herein. In at least one embodiment, registers described herein may be implemented by circuitry within a processor using any number of different techniques, such as dedicated physical registers, dynamically allocated physical registers using register renaming, combinations of dedicated and dynamically allocated physical registers, etc. In at least one embodiment, integer registers store 32-bit integer data. A register file of at least one embodiment also contains eight multimedia SIMD registers for packed data.

(465) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment portions or all of inference and/or training logic **1515** may be incorporated into EXE Block **3111** and other memory or registers shown or not shown. For example, in at least one embodiment, training and/or inferencing techniques described herein may use one or more of ALUs illustrated in EXE Block **3111**. Moreover, weight parameters may be stored in on-chip or off-chip memory and/or registers (shown or not shown) that configure ALUs of EXE Block **3111** to perform one or more machine learning algorithms, neural network architectures, use cases, or training techniques described herein.

(466) In at least one embodiment, processor **3100** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, processor **3100** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Processor **3100** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-30** and **32-43**.

(467) FIG. **32** illustrates a deep learning application processor **3200**, according to at least one

embodiment. In at least one embodiment, deep learning application processor **3200** uses instructions that, if executed by deep learning application processor **3200**, cause deep learning application processor **3200** to perform some or all of processes and techniques described throughout this disclosure. In at least one embodiment, deep learning application processor **3200** is an application-specific integrated circuit (ASIC). In at least one embodiment, application processor **3200** performs matrix multiply operations either “hard-wired” into hardware as a result of performing one or more instructions or both. In at least one embodiment, deep learning application processor **3200** includes, without limitation, processing clusters **3210(1)-3210(12)**, Inter-Chip Links (“ICLs”) **3220(1)-3220(12)**, Inter-Chip Controllers (“ICCs”) **3230(1)-3230(2)**, high bandwidth memory second generation (“HBM2”) **3240(1)-3240(4)**, memory controllers (“Mem Ctrlrs”) **3242(1)-3242(4)**, high bandwidth memory physical layer (“HBM PHY”) **3244(1)-3244(4)**, a management-controller central processing unit (“management-controller CPU”) **3250**, a Serial Peripheral Interface, Inter-Integrated Circuit, and General Purpose Input/Output block (“SPI, I2C, GPIO”) **3260**, a peripheral component interconnect express controller and direct memory access block (“PCIe Controller and DMA”) **3270**, and a sixteen-lane peripheral component interconnect express port (“PCI Express×16”) **3280**.

(468) In at least one embodiment, processing clusters **3210** may perform deep learning operations, including inference or prediction operations based on weight parameters calculated one or more training techniques, including those described herein. In at least one embodiment, each processing cluster **3210** may include, without limitation, any number and type of processors. In at least one embodiment, deep learning application processor **3200** may include any number and type of processing clusters **3200**. In at least one embodiment, Inter-Chip Links **3220** are bi-directional. In at least one embodiment, Inter-Chip Links **3220** and Inter-Chip Controllers **3230** enable multiple deep learning application processors **3200** to exchange information, including activation information resulting from performing one or more machine learning algorithms embodied in one or more neural networks. In at least one embodiment, deep learning application processor **3200** may include any number (including zero) and type of ICLs **3220** and ICCs **3230**.

(469) In at least one embodiment, HBM2s **3240** provide a total of 32 Gigabytes (GB) of memory. HBM2 **3240(i)** is associated with both memory controller **3242(i)** and HBM PHY **3244(i)**. In at least one embodiment, any number of HBM2s **3240** may provide any type and total amount of high bandwidth memory and may be associated with any number (including zero) and type of memory controllers **3242** and HBM PHYs **3244**. In at least one embodiment, SPI, I2C, GPIO **3260**, PCIe Controller and DMA **3270**, and/or PCIe **3280** may be replaced with any number and type of blocks that enable any number and type of communication standards in any technically feasible fashion.

(470) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, deep learning application processor is used to train a machine learning model, such as a neural network, to predict or infer information provided to deep learning application processor **3200**. In at least one embodiment, deep learning application processor **3200** is used to infer or predict information based on a trained machine learning model (e.g., neural network) that has been trained by another processor or system or by deep learning application processor **3200**. In at least one embodiment, processor **3200** may be used to perform one or more neural network use cases described herein.

(471) In at least one embodiment, deep learning application processor **3200** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, deep learning application processor **3200** utilizes computing resources (e.g., CPUs, ASICs, GPUs,

FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Deep learning application processor **3200** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-31** and **33-43**.

(472) FIG. **33** is a block diagram of a neuromorphic processor **3300**, according to at least one embodiment. In at least one embodiment, neuromorphic processor **3300** may receive one or more inputs from sources external to neuromorphic processor **3300**. In at least one embodiment, these inputs may be transmitted to one or more neurons **3302** within neuromorphic processor **3300**. In at least one embodiment, neurons **3302** and components thereof may be implemented using circuitry or logic, including one or more arithmetic logic units (ALUs). In at least one embodiment, neuromorphic processor **3300** may include, without limitation, thousands or millions of instances of neurons **3302**, but any suitable number of neurons **3302** may be used. In at least one embodiment, each instance of neuron **3302** may include a neuron input **3304** and a neuron output **3306**. In at least one embodiment, neurons **3302** may generate outputs that may be transmitted to inputs of other instances of neurons **3302**. For example, in at least one embodiment, neuron inputs **3304** and neuron outputs **3306** may be interconnected via synapses **3308**.

(473) In at least one embodiment, neurons **3302** and synapses **3308** may be interconnected such that neuromorphic processor **3300** operates to process or analyze information received by neuromorphic processor **3300**. In at least one embodiment, neurons **3302** may transmit an output pulse (or “fire” or “spike”) when inputs received through neuron input **3304** exceed a threshold. In at least one embodiment, neurons **3302** may sum or integrate signals received at neuron inputs **3304**. For example, in at least one embodiment, neurons **3302** may be implemented as leaky integrate-and-fire neurons, wherein if a sum (referred to as a “membrane potential”) exceeds a threshold value, neuron **3302** may generate an output (or “fire”) using a transfer function such as a sigmoid or threshold function. In at least one embodiment, a leaky integrate-and-fire neuron may sum signals received at neuron inputs **3304** into a membrane potential and may also apply a decay factor (or leak) to reduce a membrane potential. In at least one embodiment, a leaky integrate-and-fire neuron may fire if multiple input signals are received at neuron inputs **3304** rapidly enough to exceed a threshold value (i.e., before a membrane potential decays too low to fire). In at least one embodiment, neurons **3302** may be implemented using circuits or logic that receive inputs, integrate inputs into a membrane potential, and decay a membrane potential. In at least one embodiment, inputs may be averaged, or any other suitable transfer function may be used. Furthermore, in at least one embodiment, neurons **3302** may include, without limitation, comparator circuits or logic that generate an output spike at neuron output **3306** when result of applying a transfer function to neuron input **3304** exceeds a threshold. In at least one embodiment, once neuron **3302** fires, it may disregard previously received input information by, for example, resetting a membrane potential to 0 or another suitable default value. In at least one embodiment, once membrane potential is reset to 0, neuron **3302** may resume normal operation after a suitable period of time (or refractory period).

(474) In at least one embodiment, neurons **3302** may be interconnected through synapses **3308**. In at least one embodiment, synapses **3308** may operate to transmit signals from an output of a first neuron **3302** to an input of a second neuron **3302**. In at least one embodiment, neurons **3302** may transmit information over more than one instance of synapse **3308**. In at least one embodiment, one or more instances of neuron output **3306** may be connected, via an instance of synapse **3308**, to an instance of neuron input **3304** in same neuron **3302**. In at least one embodiment, an instance of neuron **3302** generating an output to be transmitted over an instance of synapse **3308** may be referred to as a “pre-synaptic neuron” with respect to that instance of synapse **3308**. In at least one embodiment, an instance of neuron **3302** receiving an input transmitted over an instance of synapse **3308** may be referred to as a “post-synaptic neuron” with respect to that instance of synapse **3308**. Because an instance of neuron **3302** may receive inputs from one or more instances of synapse

3308, and may also transmit outputs over one or more instances of synapse **3308**, a single instance of neuron **3302** may therefore be both a “pre-synaptic neuron” and “post-synaptic neuron,” with respect to various instances of synapses **3308**, in at least one embodiment.

(475) In at least one embodiment, neurons **3302** may be organized into one or more layers. Each instance of neuron **3302** may have one neuron output **3306** that may fan out through one or more synapses **3308** to one or more neuron inputs **3304**. In at least one embodiment, neuron outputs **3306** of neurons **3302** in a first layer **3310** may be connected to neuron inputs **3304** of neurons **3302** in a second layer **3312**. In at least one embodiment, layer **3310** may be referred to as a “feed-forward layer.” In at least one embodiment, each instance of neuron **3302** in an instance of first layer **3310** may fan out to each instance of neuron **3302** in second layer **3312**. In at least one embodiment, first layer **3310** may be referred to as a “fully connected feed-forward layer.” In at least one embodiment, each instance of neuron **3302** in an instance of second layer **3312** may fan out to fewer than all instances of neuron **3302** in a third layer **3314**. In at least one embodiment, second layer **3312** may be referred to as a “sparsely connected feed-forward layer.” In at least one embodiment, neurons **3302** in second layer **3312** may fan out to neurons **3302** in multiple other layers, including to neurons **3302** in (same) second layer **3312**. In at least one embodiment, second layer **3312** may be referred to as a “recurrent layer.” neuromorphic processor **3300** may include, without limitation, any suitable combination of recurrent layers and feed-forward layers, including, without limitation, both sparsely connected feed-forward layers and fully connected feed-forward layers.

(476) In at least one embodiment, neuromorphic processor **3300** may include, without limitation, a reconfigurable interconnect architecture or dedicated hard wired interconnects to connect synapse **3308** to neurons **3302**. In at least one embodiment, neuromorphic processor **3300** may include, without limitation, circuitry or logic that allows synapses to be allocated to different neurons **3302** as needed based on neural network topology and neuron fan-in/out. For example, in at least one embodiment, synapses **3308** may be connected to neurons **3302** using an interconnect fabric, such as network-on-chip, or with dedicated connections. In at least one embodiment, synapse interconnections and components thereof may be implemented using circuitry or logic.

(477) In at least one embodiment, neuromorphic processor **3300** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, neuromorphic processor **3300** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Neuromorphic processor **3300** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-32** and **34-43**.

(478) FIG. **34** is a block diagram of a processing system, according to at least one embodiment. In at least one embodiment, system **3400** includes one or more processors **3402** and one or more graphics processors **3408**, and may be a single processor desktop system, a multiprocessor workstation system, or a server system having a large number of processors **3402** or processor cores **3407**. In at least one embodiment, system **3400** is a processing platform incorporated within a system-on-a-chip (SoC) integrated circuit for use in mobile, handheld, or embedded devices.

(479) In at least one embodiment, system **3400** can include, or be incorporated within a server-based gaming platform, a game console, including a game and media console, a mobile gaming console, a handheld game console, or an online game console. In at least one embodiment, system **3400** is a mobile phone, smart phone, tablet computing device or mobile Internet device. In at least one embodiment, processing system **3400** can also include, couple with, or be integrated within a wearable device, such as a smart watch wearable device, smart eyewear device, augmented reality

device, or virtual reality device. In at least one embodiment, processing system **3400** is a television or set top box device having one or more processors **3402** and a graphical interface generated by one or more graphics processors **3408**.

(480) In at least one embodiment, one or more processors **3402** each include one or more processor cores **3407** to process instructions which, when executed, perform operations for system and user software. In at least one embodiment, each of one or more processor cores **3407** is configured to process a specific instruction set **3409**. In at least one embodiment, instruction set **3409** may facilitate Complex Instruction Set Computing (CISC), Reduced Instruction Set Computing (RISC), or computing via a Very Long Instruction Word (VLIW). In at least one embodiment, processor cores **3407** may each process a different instruction set **3409**, which may include instructions to facilitate emulation of other instruction sets. In at least one embodiment, processor core **3407** may also include other processing devices, such as a Digital Signal Processor (DSP).

(481) In at least one embodiment, processor **3402** includes cache memory **3404**. In at least one embodiment, processor **3402** can have a single internal cache or multiple levels of internal cache. In at least one embodiment, cache memory is shared among various components of processor **3402**. In at least one embodiment, processor **3402** also uses an external cache (e.g., a Level-3 (L3) cache or Last Level Cache (LLC)) (not shown), which may be shared among processor cores **3407** using known cache coherency techniques. In at least one embodiment, register file **3406** is additionally included in processor **3402** which may include different types of registers for storing different types of data (e.g., integer registers, floating point registers, status registers, and an instruction pointer register). In at least one embodiment, register file **3406** may include general-purpose registers or other registers.

(482) In at least one embodiment, one or more processor(s) **3402** are coupled with one or more interface bus(es) **3410** to transmit communication signals such as address, data, or control signals between processor **3402** and other components in system **3400**. In at least one embodiment interface bus **3410**, in one embodiment, can be a processor bus, such as a version of a Direct Media Interface (DMI) bus. In at least one embodiment, interface **3410** is not limited to a DMI bus, and may include one or more Peripheral Component Interconnect buses (e.g., PCI, PCI Express), memory busses, or other types of interface busses. In at least one embodiment processor(s) **3402** include an integrated memory controller **3416** and a platform controller hub **3430**. In at least one embodiment, memory controller **3416** facilitates communication between a memory device and other components of system **3400**, while platform controller hub (PCH) **3430** provides connections to I/O devices via a local I/O bus.

(483) In at least one embodiment, memory device **3420** can be a dynamic random access memory (DRAM) device, a static random access memory (SRAM) device, flash memory device, phase-change memory device, or some other memory device having suitable performance to serve as process memory. In at least one embodiment memory device **3420** can operate as system memory for system **3400**, to store data **3422** and instructions **3421** for use when one or more processors **3402** executes an application or process. In at least one embodiment, memory controller **3416** also couples with an optional external graphics processor **3412**, which may communicate with one or more graphics processors **3408** in processors **3402** to perform graphics and media operations. In at least one embodiment, a display device **3411** can connect to processor(s) **3402**. In at least one embodiment display device **3411** can include one or more of an internal display device, as in a mobile electronic device or a laptop device or an external display device attached via a display interface (e.g., DisplayPort, etc.). In at least one embodiment, display device **3411** can include a head mounted display (HMD) such as a stereoscopic display device for use in virtual reality (VR) applications or augmented reality (AR) applications.

(484) In at least one embodiment, platform controller hub **3430** enables peripherals to connect to memory device **3420** and processor **3402** via a high-speed I/O bus. In at least one embodiment, I/O peripherals include, but are not limited to, an audio controller **3446**, a network controller **3434**, a

firmware interface **3428**, a wireless transceiver **3426**, touch sensors **3425**, a data storage device **3424** (e.g., hard disk drive, flash memory, etc.). In at least one embodiment, data storage device **3424** can connect via a storage interface (e.g., SATA) or via a peripheral bus, such as a Peripheral Component Interconnect bus (e.g., PCI, PCI Express). In at least one embodiment, touch sensors **3425** can include touch screen sensors, pressure sensors, or fingerprint sensors. In at least one embodiment, wireless transceiver **3426** can be a Wi-Fi transceiver, a Bluetooth transceiver, or a mobile network transceiver such as a 3G, 4G, or Long Term Evolution (LTE) transceiver. In at least one embodiment, firmware interface **3428** enables communication with system firmware, and can be, for example, a unified extensible firmware interface (UEFI). In at least one embodiment, network controller **3434** can enable a network connection to a wired network. In at least one embodiment, a high-performance network controller (not shown) couples with interface bus **3410**. In at least one embodiment, audio controller **3446** is a multi-channel high definition audio controller. In at least one embodiment, system **3400** includes an optional legacy I/O controller **3440** for coupling legacy (e.g., Personal System 2 (PS/2)) devices to system. In at least one embodiment, platform controller hub **3430** can also connect to one or more Universal Serial Bus (USB) controllers **3442** connect input devices, such as keyboard and mouse **3443** combinations, a camera **3444**, or other USB input devices.

(485) In at least one embodiment, an instance of memory controller **3416** and platform controller hub **3430** may be integrated into a discreet external graphics processor, such as external graphics processor **3412**. In at least one embodiment, platform controller hub **3430** and/or memory controller **3416** may be external to one or more processor(s) **3402**. For example, in at least one embodiment, system **3400** can include an external memory controller **3416** and platform controller hub **3430**, which may be configured as a memory controller hub and peripheral controller hub within a system chipset that is in communication with processor(s) **3402**.

(486) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment portions or all of inference and/or training logic **1515** may be incorporated into graphics processor **3400**. For example, in at least one embodiment, training and/or inferencing techniques described herein may use one or more of ALUs embodied in 3D pipeline **3412**. Moreover, in at least one embodiment, inferencing and/or training operations described herein may be done using logic other than logic illustrated in FIG. **15A** or **15B**. In at least one embodiment, weight parameters may be stored in on-chip or off-chip memory and/or registers (shown or not shown) that configure ALUs of graphics processor **3400** to perform one or more machine learning algorithms, neural network architectures, use cases, or training techniques described herein.

(487) In at least one embodiment, graphics processor **3400** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, graphics processor **3400** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Graphics processor **3400** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-33** and **35-43**.

(488) FIG. **35** is a block diagram of a processor **3500** having one or more processor cores **3502A-3502N**, an integrated memory controller **3514**, and an integrated graphics processor **3508**, according to at least one embodiment. In at least one embodiment, processor **3500** can include additional cores up to and including additional core **3502N** represented by dashed lined boxes. In at least one embodiment, each of processor cores **3502A-3502N** includes one or more internal cache

units **3504A-3504N**. In at least one embodiment, each processor core also has access to one or more shared cached units **3506**.

(489) In at least one embodiment, internal cache units **3504A-3504N** and shared cache units **3506** represent a cache memory hierarchy within processor **3500**. In at least one embodiment, cache memory units **3504A-3504N** may include at least one level of instruction and data cache within each processor core and one or more levels of shared mid-level cache, such as a Level 2 (L2), Level 3 (L3), Level 4 (L4), or other levels of cache, where a highest level of cache before external memory is classified as an LLC. In at least one embodiment, cache coherency logic maintains coherency between various cache units **3506** and **3504A-3504N**.

(490) In at least one embodiment, processor **3500** may also include a set of one or more bus controller units **3516** and a system agent core **3510**. In at least one embodiment, one or more bus controller units **3516** manage a set of peripheral buses, such as one or more PCI or PCI express busses. In at least one embodiment, system agent core **3510** provides management functionality for various processor components. In at least one embodiment, system agent core **3510** includes one or more integrated memory controllers **3514** to manage access to various external memory devices (not shown).

(491) In at least one embodiment, one or more of processor cores **3502A-3502N** include support for simultaneous multi-threading. In at least one embodiment, system agent core **3510** includes components for coordinating and operating cores **3502A-3502N** during multi-threaded processing. In at least one embodiment, system agent core **3510** may additionally include a power control unit (PCU), which includes logic and components to regulate one or more power states of processor cores **3502A-3502N** and graphics processor **3508**.

(492) In at least one embodiment, processor **3500** additionally includes graphics processor **3508** to execute graphics processing operations. In at least one embodiment, graphics processor **3508** couples with shared cache units **3506**, and system agent core **3510**, including one or more integrated memory controllers **3514**. In at least one embodiment, system agent core **3510** also includes a display controller **3511** to drive graphics processor output to one or more coupled displays. In at least one embodiment, display controller **3511** may also be a separate module coupled with graphics processor **3508** via at least one interconnect, or may be integrated within graphics processor **3508**.

(493) In at least one embodiment, a ring based interconnect unit **3512** is used to couple internal components of processor **3500**. In at least one embodiment, an alternative interconnect unit may be used, such as a point-to-point interconnect, a switched interconnect, or other techniques. In at least one embodiment, graphics processor **3508** couples with ring interconnect **3512** via an I/O link **3513**.

(494) In at least one embodiment, I/O link **3513** represents at least one of multiple varieties of I/O interconnects, including an on package I/O interconnect which facilitates communication between various processor components and a high-performance embedded memory module **3518**, such as an eDRAM module. In at least one embodiment, each of processor cores **3502A-3502N** and graphics processor **3508** use embedded memory modules **3518** as a shared Last Level Cache.

(495) In at least one embodiment, processor cores **3502A-3502N** are homogenous cores executing a common instruction set architecture. In at least one embodiment, processor cores **3502A-3502N** are heterogeneous in terms of instruction set architecture (ISA), where one or more of processor cores **3502A-3502N** execute a common instruction set, while one or more other cores of processor cores **3502A-35-02N** executes a subset of a common instruction set or a different instruction set. In at least one embodiment, processor cores **3502A-3502N** are heterogeneous in terms of microarchitecture, where one or more cores having a relatively higher power consumption couple with one or more power cores having a lower power consumption. In at least one embodiment, processor **3500** can be implemented on one or more chips or as an SoC integrated circuit.

(496) Inference and/or training logic **1515** are used to perform inferencing and/or training

operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment portions or all of inference and/or training logic **1515** may be incorporated into graphics processor **3510**. For example, in at least one embodiment, training and/or inferencing techniques described herein may use one or more of ALUs embodied in 3D pipeline **3412**, graphics core(s) **3515A**, shared function logic **3516**, graphics core(s) **3515B**, shared function logic **3520**, or other logic in FIG. **35**. Moreover, in at least one embodiment, inferencing and/or training operations described herein may be done using logic other than logic illustrated in FIG. **15A** or **15B**. In at least one embodiment, weight parameters may be stored in on-chip or off-chip memory and/or registers (shown or not shown) that configure ALUs of graphics processor **3510** to perform one or more machine learning algorithms, neural network architectures, use cases, or training techniques described herein.

(497) In at least one embodiment, processor **3500** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, processor **3500** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Processor **3500** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-34** and **36-43**.

(498) FIG. **36** is a block diagram of a graphics processor **3600**, which may be a discrete graphics processing unit, or may be a graphics processor integrated with a plurality of processing cores. In at least one embodiment, graphics processor **3600** communicates via a memory mapped I/O interface to registers on graphics processor **3600** and with commands placed into memory. In at least one embodiment, graphics processor **3600** includes a memory interface **3614** to access memory. In at least one embodiment, memory interface **3614** is an interface to local memory, one or more internal caches, one or more shared external caches, and/or to system memory.

(499) In at least one embodiment, graphics processor **3600** also includes a display controller **3602** to drive display output data to a display device **3620**. In at least one embodiment, display controller **3602** includes hardware for one or more overlay planes for display device **3620** and composition of multiple layers of video or user interface elements. In at least one embodiment, display device **3620** can be an internal or external display device. In at least one embodiment, display device **3620** is a head mounted display device, such as a virtual reality (VR) display device or an augmented reality (AR) display device. In at least one embodiment, graphics processor **3600** includes a video codec engine **3606** to encode, decode, or transcode media to, from, or between one or more media encoding formats, including, but not limited to Moving Picture Experts Group (MPEG) formats such as MPEG-2, Advanced Video Coding (AVC) formats such as H.264/MPEG-4 AVC, as well as the Society of Motion Picture & Television Engineers (SMPTE) 421M/VC-1, and Joint Photographic Experts Group (JPEG) formats such as JPEG, and Motion JPEG (MJPEG) formats.

(500) In at least one embodiment, graphics processor **3600** includes a block image transfer (BLIT) engine **3604** to perform two-dimensional (2D) rasterizer operations including, for example, bit-boundary block transfers. However, in at least one embodiment, 2D graphics operations are performed using one or more components of graphics processing engine (GPE) **3610**. In at least one embodiment, GPE **3610** is a compute engine for performing graphics operations, including three-dimensional (3D) graphics operations and media operations.

(501) In at least one embodiment, GPE **3610** includes a 3D pipeline **3612** for performing 3D operations, such as rendering three-dimensional images and scenes using processing functions that act upon 3D primitive shapes (e.g., rectangle, triangle, etc.). 3D pipeline **3612** includes programmable and fixed function elements that perform various tasks and/or spawn execution

threads to a 3D/Media sub-system **3615**. While 3D pipeline **3612** can be used to perform media operations, in at least one embodiment, GPE **3610** also includes a media pipeline **3616** that is used to perform media operations, such as video post-processing and image enhancement.

(502) In at least one embodiment, media pipeline **3616** includes fixed function or programmable logic units to perform one or more specialized media operations, such as video decode acceleration, video de-interlacing, and video encode acceleration in place of, or on behalf of video codec engine **3606**. In at least one embodiment, media pipeline **3616** additionally includes a thread spawning unit to spawn threads for execution on 3D/Media sub-system **3615**. In at least one embodiment, spawned threads perform computations for media operations on one or more graphics execution units included in 3D/Media sub-system **3615**.

(503) In at least one embodiment, 3D/Media subsystem **3615** includes logic for executing threads spawned by 3D pipeline **3612** and media pipeline **3616**. In at least one embodiment, 3D pipeline **3612** and media pipeline **3616** send thread execution requests to 3D/Media subsystem **3615**, which includes thread dispatch logic for arbitrating and dispatching various requests to available thread execution resources. In at least one embodiment, execution resources include an array of graphics execution units to process 3D and media threads. In at least one embodiment, 3D/Media subsystem **3615** includes one or more internal caches for thread instructions and data. In at least one embodiment, subsystem **3615** also includes shared memory, including registers and addressable memory, to share data between threads and to store output data.

(504) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment portions or all of inference and/or training logic **1515** may be incorporated into graphics processor **3600**. For example, in at least one embodiment, training and/or inferencing techniques described herein may use one or more of ALUs embodied in 3D pipeline **3612**.

Moreover, in at least one embodiment, inferencing and/or training operations described herein may be done using logic other than logic illustrated in FIG. **15A** or **15B**. In at least one embodiment, weight parameters may be stored in on-chip or off-chip memory and/or registers (shown or not shown) that configure ALUs of graphics processor **3600** to perform one or more machine learning algorithms, neural network architectures, use cases, or training techniques described herein.

(505) In at least one embodiment, graphics processor **3600** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, graphics processor **3600** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Graphics processor **3600** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-35** and **37-43**.

(506) FIG. **37** is a block diagram of a graphics processing engine **3710** of a graphics processor in accordance with at least one embodiment. In at least one embodiment, graphics processing engine (GPE) **3710** is a version of GPE **3610** shown in FIG. **36**. In at least one embodiment, media pipeline **3616** is optional and may not be explicitly included within GPE **3710**. In at least one embodiment, a separate media and/or image processor is coupled to GPE **3710**.

(507) In at least one embodiment, GPE **3710** is coupled to or includes a command streamer **3703**, which provides a command stream to 3D pipeline **3612** and/or media pipelines **3616**. In at least one embodiment, command streamer **3703** is coupled to memory, which can be system memory, or one or more of internal cache memory and shared cache memory. In at least one embodiment, command streamer **3703** receives commands from memory and sends commands to 3D pipeline

3612 and/or media pipeline **3616**. In at least one embodiment, commands are instructions, primitives, or micro-operations fetched from a ring buffer, which stores commands for 3D pipeline **3612** and media pipeline **3616**. In at least one embodiment, a ring buffer can additionally include batch command buffers storing batches of multiple commands. In at least one embodiment, commands for 3D pipeline **3612** can also include references to data stored in memory, such as but not limited to vertex and geometry data for 3D pipeline **3612** and/or image data and memory objects for media pipeline **3616**. In at least one embodiment, 3D pipeline **3612** and media pipeline **3616** process commands and data by performing operations or by dispatching one or more execution threads to a graphics core array **3714**. In at least one embodiment graphics core array **3714** includes one or more blocks of graphics cores (e.g., graphics core(s) **3715A**, graphics core(s) **3715B**), each block including one or more graphics cores. In at least one embodiment, each graphics core includes a set of graphics execution resources that includes general-purpose and graphics specific execution logic to perform graphics and compute operations, as well as fixed function texture processing and/or machine learning and artificial intelligence acceleration logic, including inference and/or training logic **1515** in FIG. 15A and FIG. 15B.

(508) In at least one embodiment, 3D pipeline **3612** includes fixed function and programmable logic to process one or more shader programs, such as vertex shaders, geometry shaders, pixel shaders, fragment shaders, compute shaders, or other shader programs, by processing instructions and dispatching execution threads to graphics core array **3714**. In at least one embodiment, graphics core array **3714** provides a unified block of execution resources for use in processing shader programs. In at least one embodiment, multi-purpose execution logic (e.g., execution units) within graphics core(s) **3715A-3715B** of graphic core array **3714** includes support for various 3D API shader languages and can execute multiple simultaneous execution threads associated with multiple shaders.

(509) In at least one embodiment, graphics core array **3714** also includes execution logic to perform media functions, such as video and/or image processing. In at least one embodiment, execution units additionally include general-purpose logic that is programmable to perform parallel general-purpose computational operations, in addition to graphics processing operations.

(510) In at least one embodiment, output data generated by threads executing on graphics core array **3714** can output data to memory in a unified return buffer (URB) **3718**. URB **3718** can store data for multiple threads. In at least one embodiment, URB **3718** may be used to send data between different threads executing on graphics core array **3714**. In at least one embodiment, URB **3718** may additionally be used for synchronization between threads on graphics core array **3714** and fixed function logic within shared function logic **3720**.

(511) In at least one embodiment, graphics core array **3714** is scalable, such that graphics core array **3714** includes a variable number of graphics cores, each having a variable number of execution units based on a target power and performance level of GPE **3710**. In at least one embodiment, execution resources are dynamically scalable, such that execution resources may be enabled or disabled as needed.

(512) In at least one embodiment, graphics core array **3714** is coupled to shared function logic **3720** that includes multiple resources that are shared between graphics cores in graphics core array **3714**. In at least one embodiment, shared functions performed by shared function logic **3720** are embodied in hardware logic units that provide specialized supplemental functionality to graphics core array **3714**. In at least one embodiment, shared function logic **3720** includes but is not limited to sampler **3721**, math **3722**, and inter-thread communication (ITC) **3723** logic. In at least one embodiment, one or more cache(s) **3725** are included in or couple to shared function logic **3720**.

(513) In at least one embodiment, a shared function is used if demand for a specialized function is insufficient for inclusion within graphics core array **3714**. In at least one embodiment, a single instantiation of a specialized function is used in shared function logic **3720** and shared among other execution resources within graphics core array **3714**. In at least one embodiment, specific shared

functions within shared function logic **3720** that are used extensively by graphics core array **3714** may be included within shared function logic **3716** within graphics core array **3714**. In at least one embodiment, shared function logic **3716** within graphics core array **3714** can include some or all logic within shared function logic **3720**. In at least one embodiment, all logic elements within shared function logic **3720** may be duplicated within shared function logic **3716** of graphics core array **3714**. In at least one embodiment, shared function logic **3720** is excluded in favor of shared function logic **3716** within graphics core array **3714**.

(514) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment portions or all of inference and/or training logic **1515** may be incorporated into graphics processor **3710**. For example, in at least one embodiment, training and/or inferencing techniques described herein may use one or more of ALUs embodied in 3D pipeline **3612**, graphics core(s) **3715A**, shared function logic **3716**, graphics core(s) **3715B**, shared function logic **3720**, or other logic in FIG. **37**. Moreover, in at least one embodiment, inferencing and/or training operations described herein may be done using logic other than logic illustrated in FIG. **15A** or **15B**. In at least one embodiment, weight parameters may be stored in on-chip or off-chip memory and/or registers (shown or not shown) that configure ALUs of graphics processor **3710** to perform one or more machine learning algorithms, neural network architectures, use cases, or training techniques described herein.

(515) In at least one embodiment, graphics processing engine **3710** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, graphics processing engine **3710** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Graphics processing engine **3710** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-36** and **38-43**.

(516) FIG. **38** is a block diagram of hardware logic of a graphics processor core **3800**, according to at least one embodiment described herein. In at least one embodiment, graphics processor core **3800** is included within a graphics core array. In at least one embodiment, graphics processor core **3800**, sometimes referred to as a core slice, can be one or multiple graphics cores within a modular graphics processor. In at least one embodiment, graphics processor core **3800** is exemplary of one graphics core slice, and a graphics processor as described herein may include multiple graphics core slices based on target power and performance envelopes. In at least one embodiment, each graphics core **3800** can include a fixed function block **3830** coupled with multiple sub-cores **3801A-3801F**, also referred to as sub-slices, that include modular blocks of general-purpose and fixed function logic.

(517) In at least one embodiment, fixed function block **3830** includes a geometry/fixed function pipeline **3836** that can be shared by all sub-cores in graphics processor **3800**, for example, in lower performance and/or lower power graphics processor implementations. In at least one embodiment, geometry/fixed function pipeline **3836** includes a 3D fixed function pipeline, a video front-end unit, a thread spawner and thread dispatcher, and a unified return buffer manager, which manages unified return buffers.

(518) In at least one embodiment fixed function block **3830** also includes a graphics SoC interface **3837**, a graphics microcontroller **3838**, and a media pipeline **3839**. Graphics SoC interface **3837** provides an interface between graphics core **3800** and other processor cores within a system on a chip integrated circuit. In at least one embodiment, graphics microcontroller **3838** is a

programmable sub-processor that is configurable to manage various functions of graphics processor **3800**, including thread dispatch, scheduling, and pre-emption. In at least one embodiment, media pipeline **3839** includes logic to facilitate decoding, encoding, pre-processing, and/or post-processing of multimedia data, including image and video data. In at least one embodiment, media pipeline **3839** implement media operations via requests to compute or sampling logic within sub-cores **3801-3801F**.

(519) In at least one embodiment, SoC interface **3837** enables graphics core **3800** to communicate with general-purpose application processor cores (e.g., CPUs) and/or other components within an SoC, including memory hierarchy elements such as a shared last level cache memory, system RAM, and/or embedded on-chip or on-package DRAM. In at least one embodiment, SoC interface **3837** can also enable communication with fixed function devices within an SoC, such as camera imaging pipelines, and enables use of and/or implements global memory atomics that may be shared between graphics core **3800** and CPUs within an SoC. In at least one embodiment, SoC interface **3837** can also implement power management controls for graphics core **3800** and enable an interface between a clock domain of graphic core **3800** and other clock domains within an SoC. In at least one embodiment, SoC interface **3837** enables receipt of command buffers from a command streamer and global thread dispatcher that are configured to provide commands and instructions to each of one or more graphics cores within a graphics processor. In at least one embodiment, commands and instructions can be dispatched to media pipeline **3839**, when media operations are to be performed, or a geometry and fixed function pipeline (e.g., geometry and fixed function pipeline **3836**, geometry and fixed function pipeline **3814**) when graphics processing operations are to be performed.

(520) In at least one embodiment, graphics microcontroller **3838** can be configured to perform various scheduling and management tasks for graphics core **3800**. In at least one embodiment, graphics microcontroller **3838** can perform graphics and/or compute workload scheduling on various graphics parallel engines within execution unit (EU) arrays **3802A-3802F**, **3804A-3804F** within sub-cores **3801A-3801F**. In at least one embodiment, host software executing on a CPU core of an SoC including graphics core **3800** can submit workloads one of multiple graphic processor doorbells, which invokes a scheduling operation on an appropriate graphics engine. In at least one embodiment, scheduling operations include determining which workload to run next, submitting a workload to a command streamer, pre-empting existing workloads running on an engine, monitoring progress of a workload, and notifying host software when a workload is complete. In at least one embodiment, graphics microcontroller **3838** can also facilitate low-power or idle states for graphics core **3800**, providing graphics core **3800** with an ability to save and restore registers within graphics core **3800** across low-power state transitions independently from an operating system and/or graphics driver software on a system.

(521) In at least one embodiment, graphics core **3800** may have greater than or fewer than illustrated sub-cores **3801A-3801F**, up to N modular sub-cores. For each set of N sub-cores, in at least one embodiment, graphics core **3800** can also include shared function logic **3810**, shared and/or cache memory **3812**, a geometry/fixed function pipeline **3814**, as well as additional fixed function logic **3816** to accelerate various graphics and compute processing operations. In at least one embodiment, shared function logic **3810** can include logic units (e.g., sampler, math, and/or inter-thread communication logic) that can be shared by each N sub-cores within graphics core **3800**. Shared and/or cache memory **3812** can be a last-level cache for N sub-cores **3801A-3801F** within graphics core **3800** and can also serve as shared memory that is accessible by multiple sub-cores. In at least one embodiment, geometry/fixed function pipeline **3814** can be included instead of geometry/fixed function pipeline **3836** within fixed function block **3830** and can include same or similar logic units.

(522) In at least one embodiment, graphics core **3800** includes additional fixed function logic **3816** that can include various fixed function acceleration logic for use by graphics core **3800**. In at least

one embodiment, additional fixed function logic **3816** includes an additional geometry pipeline for use in position only shading. In position-only shading, at least two geometry pipelines exist, whereas in a full geometry pipeline within geometry/fixed function pipeline **3816**, **3836**, and a cull pipeline, which is an additional geometry pipeline which may be included within additional fixed function logic **3816**. In at least one embodiment, cull pipeline is a trimmed down version of a full geometry pipeline. In at least one embodiment, a full pipeline and a cull pipeline can execute different instances of an application, each instance having a separate context. In at least one embodiment, position only shading can hide long cull runs of discarded triangles, enabling shading to be completed earlier in some instances. For example, in at least one embodiment, cull pipeline logic within additional fixed function logic **3816** can execute position shaders in parallel with a main application and generally generates critical results faster than a full pipeline, as cull pipeline fetches and shades position attribute of vertices, without performing rasterization and rendering of pixels to a frame buffer. In at least one embodiment, cull pipeline can use generated critical results to compute visibility information for all triangles without regard to whether those triangles are culled. In at least one embodiment, full pipeline (which in this instance may be referred to as a replay pipeline) can consume visibility information to skip culled triangles to shade only visible triangles that are finally passed to a rasterization phase.

(523) In at least one embodiment, additional fixed function logic **3816** can also include machine-learning acceleration logic, such as fixed function matrix multiplication logic, for implementations including optimizations for machine learning training or inferencing.

(524) In at least one embodiment, within each graphics sub-core **3801A-3801F** includes a set of execution resources that may be used to perform graphics, media, and compute operations in response to requests by graphics pipeline, media pipeline, or shader programs. In at least one embodiment, graphics sub-cores **3801A-3801F** include multiple EU arrays **3802A-3802F**, **3804A-3804F**, thread dispatch and inter-thread communication (TD/IC) logic **3803A-3803F**, a 3D (e.g., texture) sampler **3805A-3805F**, a media sampler **3806A-3806F**, a shader processor **3807A-3807F**, and shared local memory (SLM) **3808A-3808F**. EU arrays **3802A-3802F**, **3804A-3804F** each include multiple execution units, which are general-purpose graphics processing units capable of performing floating-point and integer/fixed-point logic operations in service of a graphics, media, or compute operation, including graphics, media, or compute shader programs. In at least one embodiment, TD/IC logic **3803A-3803F** performs local thread dispatch and thread control operations for execution units within a sub-core and facilitate communication between threads executing on execution units of a sub-core. In at least one embodiment, 3D sampler **3805A-3805F** can read texture or other 3D graphics related data into memory. In at least one embodiment, 3D sampler can read texture data differently based on a configured sample state and texture format associated with a given texture. In at least one embodiment, media sampler **3806A-3806F** can perform similar read operations based on a type and format associated with media data. In at least one embodiment, each graphics sub-core **3801A-3801F** can alternately include a unified 3D and media sampler. In at least one embodiment, threads executing on execution units within each of sub-cores **3801A-3801F** can make use of shared local memory **3808A-3808F** within each sub-core, to enable threads executing within a thread group to execute using a common pool of on-chip memory.

(525) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, portions or all of inference and/or training logic **1515** may be incorporated into graphics processor **3810**. For example, in at least one embodiment, training and/or inferencing techniques described herein may use one or more of ALUs embodied in 3D pipeline **3810**, graphics microcontroller **3838**, geometry & fixed function pipeline **3814** and **3836**, or other logic in FIG. **35**. Moreover, in at least one embodiment, inferencing and/or training operations described herein may

be done using logic other than logic illustrated in FIG. 15A or 15B. In at least one embodiment, weight parameters may be stored in on-chip or off-chip memory and/or registers (shown or not shown) that configure ALUs of graphics processor **3800** to perform one or more machine learning algorithms, neural network architectures, use cases, or training techniques described herein. (526) In at least one embodiment, graphics processor **3800** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, graphics processor **3800** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Graphics processor **3800** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-37** and **39-43**.

(527) FIGS. **39A-39B** illustrate thread execution logic **3900** including an array of processing elements of a graphics processor core according to at least one embodiment. FIG. **39A** illustrates at least one embodiment, in which thread execution logic **3900** is used. FIG. **39B** illustrates exemplary internal details of an execution unit, according to at least one embodiment.

(528) As illustrated in FIG. **39A**, in at least one embodiment, thread execution logic **3900** includes a shader processor **3902**, a thread dispatcher **3904**, instruction cache **3906**, a scalable execution unit array including a plurality of execution units **3908A-3908N**, a sampler **3910**, a data cache **3912**, and a data port **3914**. In at least one embodiment a scalable execution unit array can dynamically scale by enabling or disabling one or more execution units (e.g., any of execution unit **3908A**, **3908B**, **3908C**, **3908D**, through **3908N-1** and **3908N**) based on computational requirements of a workload, for example. In at least one embodiment, scalable execution units are interconnected via an interconnect fabric that links to each of execution unit. In at least one embodiment, thread execution logic **3900** includes one or more connections to memory, such as system memory or cache memory, through one or more of instruction cache **3906**, data port **3914**, sampler **3910**, and execution units **3908A-3908N**. In at least one embodiment, each execution unit (e.g., **3908A**) is a stand-alone programmable general-purpose computational unit that is capable of executing multiple simultaneous hardware threads while processing multiple data elements in parallel for each thread. In at least one embodiment, array of execution units **3908A-3908N** is scalable to include any number individual execution units.

(529) In at least one embodiment, execution units **3908A-3908N** are primarily used to execute shader programs. In at least one embodiment, shader processor **3902** can process various shader programs and dispatch execution threads associated with shader programs via a thread dispatcher **3904**. In at least one embodiment, thread dispatcher **3904** includes logic to arbitrate thread initiation requests from graphics and media pipelines and instantiate requested threads on one or more execution units in execution units **3908A-3908N**. For example, in at least one embodiment, a geometry pipeline can dispatch vertex, tessellation, or geometry shaders to thread execution logic for processing. In at least one embodiment, thread dispatcher **3904** can also process runtime thread spawning requests from executing shader programs.

(530) In at least one embodiment, execution units **3908A-3908N** support an instruction set that includes native support for many standard 3D graphics shader instructions, such that shader programs from graphics libraries (e.g., Direct 3D and OpenGL) are executed with a minimal translation. In at least one embodiment, execution units support vertex and geometry processing (e.g., vertex programs, geometry programs, vertex shaders), pixel processing (e.g., pixel shaders, fragment shaders) and general-purpose processing (e.g., compute and media shaders). In at least one embodiment, each of execution units **3908A-3908N**, which include one or more arithmetic logic units (ALUs), is capable of multi-issue single instruction multiple data (SIMD) execution and

multi-threaded operation enables an efficient execution environment despite higher latency memory accesses. In at least one embodiment, each hardware thread within each execution unit has a dedicated high-bandwidth register file and associated independent thread-state. In at least one embodiment, execution is multi-issue per clock to pipelines capable of integer, single and double precision floating point operations, SIMD branch capability, logical operations, transcendental operations, and other miscellaneous operations. In at least one embodiment, while waiting for data from memory or one of shared functions, dependency logic within execution units **3908A-3908N** causes a waiting thread to sleep until requested data has been returned. In at least one embodiment, while a waiting thread is sleeping, hardware resources may be devoted to processing other threads. For example, in at least one embodiment, during a delay associated with a vertex shader operation, an execution unit can perform operations for a pixel shader, fragment shader, or another type of shader program, including a different vertex shader.

(531) In at least one embodiment, each execution unit in execution units **3908A-3908N** operates on arrays of data elements. In at least one embodiment, a number of data elements is “execution size,” or number of channels for an instruction. In at least one embodiment, an execution channel is a logical unit of execution for data element access, masking, and flow control within instructions. In at least one embodiment, a number of channels may be independent of a number of physical Arithmetic Logic Units (ALUs) or Floating Point Units (FPUs) for a particular graphics processor. In at least one embodiment, execution units **3908A-3908N** support integer and floating-point data types.

(532) In at least one embodiment, an execution unit instruction set includes SIMD instructions. In at least one embodiment, various data elements can be stored as a packed data type in a register and execution unit will process various elements based on data size of elements. For example, in at least one embodiment, when operating on a 256-bit wide vector, 256 bits of a vector are stored in a register and an execution unit operates on a vector as four separate 64-bit packed data elements (Quad-Word (QW) size data elements), eight separate 32-bit packed data elements (Double Word (DW) size data elements), sixteen separate 16-bit packed data elements (Word (W) size data elements), or thirty-two separate 8-bit data elements (byte (B) size data elements). However, in at least one embodiment, different vector widths and register sizes are possible.

(533) In at least one embodiment, one or more execution units can be combined into a fused execution unit **3909A-3909N** having thread control logic (**3907A-3907N**) that is common to fused EUs. In at least one embodiment, multiple EUs can be fused into an EU group. In at least one embodiment, each EU in fused EU group can be configured to execute a separate SIMD hardware thread. The number of EUs in a fused EU group can vary according to various embodiments. In at least one embodiment, various SIMD widths can be performed per-EU, including but not limited to SIMD8, SIMD16, and SIMD32. In at least one embodiment, each fused graphics execution unit **3909A-3909N** includes at least two execution units. For example, in at least one embodiment, fused execution unit **3909A** includes a first EU **3908A**, second EU **3908B**, and thread control logic **3907A** that is common to first EU **3908A** and second EU **3908B**. In at least one embodiment, thread control logic **3907A** controls threads executed on fused graphics execution unit **3909A**, allowing each EU within fused execution units **3909A-3909N** to execute using a common instruction pointer register.

(534) In at least one embodiment, one or more internal instruction caches (e.g., **3906**) are included in thread execution logic **3900** to cache thread instructions for execution units. In at least one embodiment, one or more data caches (e.g., **3912**) are included to cache thread data during thread execution. In at least one embodiment, a sampler **3910** is included to provide texture sampling for 3D operations and media sampling for media operations. In at least one embodiment, sampler **3910** includes specialized texture or media sampling functionality to process texture or media data during sampling process before providing sampled data to an execution unit.

(535) During execution, in at least one embodiment, graphics and media pipelines send thread

initiation requests to thread execution logic **3900** via thread spawning and dispatch logic. In at least one embodiment, once a group of geometric objects has been processed and rasterized into pixel data, pixel processor logic (e.g., pixel shader logic, fragment shader logic, etc.) within shader processor **3902** is invoked to further compute output information and cause results to be written to output surfaces (e.g., color buffers, depth buffers, stencil buffers, etc.). In at least one embodiment, a pixel shader or fragment shader calculates values of various vertex attributes that are to be interpolated across a rasterized object. In at least one embodiment, pixel processor logic within shader processor **3902** then executes an application programming interface (API)-supplied pixel or fragment shader program. In at least one embodiment, to execute a shader program, shader processor **3902** dispatches threads to an execution unit (e.g., **3908A**) via thread dispatcher **3904**. In at least one embodiment, shader processor **3902** uses texture sampling logic in sampler **3910** to access texture data in texture maps stored in memory. In at least one embodiment, arithmetic operations on texture data and input geometry data compute pixel color data for each geometric fragment, or discards one or more pixels from further processing.

(536) In at least one embodiment, data port **3914** provides a memory access mechanism for thread execution logic **3900** to output processed data to memory for further processing on a graphics processor output pipeline. In at least one embodiment, data port **3914** includes or couples to one or more cache memories (e.g., data cache **3912**) to cache data for memory access via a data port.

(537) As illustrated in FIG. **39B**, in at least one embodiment, a graphics execution unit **3908** can include an instruction fetch unit **3937**, a general register file array (GRF) **3924**, an architectural register file array (ARF) **3926**, a thread arbiter **3922**, a send unit **3930**, a branch unit **3932**, a set of SIMD floating point units (FPUs) **3934**, and In at least one embodiment a set of dedicated integer SIMD ALUs **3935**. In at least one embodiment, GRF **3924** and ARF **3926** includes a set of general register files and architecture register files associated with each simultaneous hardware thread that may be active in graphics execution unit **3908**. In at least one embodiment, per thread architectural state is maintained in ARF **3926**, while data used during thread execution is stored in GRF **3924**. In at least one embodiment, execution state of each thread, including instruction pointers for each thread, can be held in thread-specific registers in ARF **3926**.

(538) In at least one embodiment, graphics execution unit **3908** has an architecture that is a combination of Simultaneous Multi-Threading (SMT) and fine-grained Interleaved Multi-Threading (IMT). In at least one embodiment, architecture has a modular configuration that can be fine-tuned at design time based on a target number of simultaneous threads and number of registers per execution unit, where execution unit resources are divided across logic used to execute multiple simultaneous threads.

(539) In at least one embodiment, graphics execution unit **3908** can co-issue multiple instructions, which may each be different instructions. In at least one embodiment, thread arbiter **3922** of graphics execution unit thread **3908** can dispatch instructions to one of send unit **3930**, branch unit **3942**, or SIMD FPU(s) **3934** for execution. In at least one embodiment, each execution thread can access **128** general-purpose registers within GRF **3924**, where each register can store 32 bytes, accessible as a SIMD 8-element vector of 32-bit data elements. In at least one embodiment, each execution unit thread has access to 4 Kbytes within GRF **3924**, although embodiments are not so limited, and greater or fewer register resources may be provided in other embodiments. In at least one embodiment, up to seven threads can execute simultaneously, although a number of threads per execution unit can also vary according to embodiments. In at least one embodiment, in which seven threads may access 4 Kbytes, GRF **3924** can store a total of 28 Kbytes. In at least one embodiment, flexible addressing modes can permit registers to be addressed together to build effectively wider registers or to represent strided rectangular block data structures.

(540) In at least one embodiment, memory operations, sampler operations, and other longer-latency system communications are dispatched via “send” instructions that are executed by message passing send unit **3930**. In at least one embodiment, branch instructions are dispatched to a

dedicated branch unit **3932** to facilitate SIMD divergence and eventual convergence.

(541) In at least one embodiment graphics execution unit **3908** includes one or more SIMD floating point units (FPU(s)) **3934** to perform floating-point operations. In at least one embodiment, FPU(s) **3934** also support integer computation. In at least one embodiment FPU(s) **3934** can SIMD execute up to M number of 32-bit floating-point (or integer) operations, or SIMD execute up to 2M 16-bit integer or 16-bit floating-point operations. In at least one embodiment, at least one of FPU(s) provides extended math capability to support high-throughput transcendental math functions and double precision 64-bit floating-point. In at least one embodiment, a set of 8-bit integer SIMD ALUs **3935** are also present, and may be specifically optimized to perform operations associated with machine learning computations.

(542) In at least one embodiment, arrays of multiple instances of graphics execution unit **3908** can be instantiated in a graphics sub-core grouping (e.g., a sub-slice). In at least one embodiment execution unit **3908** can execute instructions across a plurality of execution channels. In at least one embodiment, each thread executed on graphics execution unit **3908** is executed on a different channel.

(543) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, portions or all of inference and/or training logic **1515** may be incorporated into execution logic **3900**. Moreover, in at least one embodiment, inferencing and/or training operations described herein may be done using logic other than logic illustrated in FIG. **15A** or **15B**. In at least one embodiment, weight parameters may be stored in on-chip or off-chip memory and/or registers (shown or not shown) that configure ALUs of execution logic **3900** to perform one or more machine learning algorithms, neural network architectures, use cases, or training techniques described herein.

(544) In at least one embodiment, execution logic **3900** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, execution logic **3900** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. Execution logic **3900** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-38** and **40-43**.

(545) FIG. **40** illustrates a parallel processing unit (“PPU”) **4000**, according to at least one embodiment. In at least one embodiment, PPU **4000** is configured with machine-readable code that, if executed by PPU **4000**, causes PPU **4000** to perform some or all of processes and techniques described throughout this disclosure. In at least one embodiment, PPU **4000** is a multi-threaded processor that is implemented on one or more integrated circuit devices and that utilizes multithreading as a latency-hiding technique designed to process computer-readable instructions (also referred to as machine-readable instructions or simply instructions) on multiple threads in parallel. In at least one embodiment, a thread refers to a thread of execution and is an instantiation of a set of instructions configured to be executed by PPU **4000**. In at least one embodiment, PPU **4000** is a graphics processing unit (“GPU”) configured to implement a graphics rendering pipeline for processing three-dimensional (“3D”) graphics data in order to generate two-dimensional (“2D”) image data for display on a display device such as a liquid crystal display (“LCD”) device. In at least one embodiment, PPU **4000** is utilized to perform computations such as linear algebra operations and machine-learning operations. FIG. **40** illustrates an example parallel processor for illustrative purposes only and should be construed as a non-limiting example of processor

architectures contemplated within scope of this disclosure and that any suitable processor may be employed to supplement and/or substitute for same.

(546) In at least one embodiment, one or more PPUs **4000** are configured to accelerate High Performance Computing (“HPC”), data center, and machine learning applications. In at least one embodiment, PPU **4000** is configured to accelerate deep learning systems and applications including following non-limiting examples: autonomous vehicle platforms, deep learning, high-accuracy speech, image, text recognition systems, intelligent video analytics, molecular simulations, drug discovery, disease diagnosis, weather forecasting, big data analytics, astronomy, molecular dynamics simulation, financial modeling, robotics, factory automation, real-time language translation, online search optimizations, and personalized user recommendations, and more.

(547) In at least one embodiment, PPU **4000** includes, without limitation, an Input/Output (“I/O”) unit **4006**, a front-end unit **4010**, a scheduler unit **4012**, a work distribution unit **4014**, a hub **4016**, a crossbar (“Xbar”) **4020**, one or more general processing clusters (“GPCs”) **4018**, and one or more partition units (“memory partition units”) **4022**. In at least one embodiment, PPU **4000** is connected to a host processor or other PPUs **4000** via one or more high-speed GPU interconnects (“GPU interconnects”) **4008**. In at least one embodiment, PPU **4000** is connected to a host processor or other peripheral devices via an interconnect **4002**. In at least one embodiment, PPU **4000** is connected to a local memory comprising one or more memory devices (“memory”) **4004**. In at least one embodiment, memory devices **4004** include, without limitation, one or more dynamic random access memory (“DRAM”) devices. In at least one embodiment, one or more DRAM devices are configured and/or configurable as high-bandwidth memory (“HBM”) subsystems, with multiple DRAM dies stacked within each device.

(548) In at least one embodiment, high-speed GPU interconnect **4008** may refer to a wire-based multi-lane communications link that is used by systems to scale and include one or more PPUs **4000** combined with one or more central processing units (“CPUs”), supports cache coherence between PPUs **4000** and CPUs, and CPU mastering. In at least one embodiment, data and/or commands are transmitted by high-speed GPU interconnect **4008** through hub **4016** to/from other units of PPU **4000** such as one or more copy engines, video encoders, video decoders, power management units, and other components which may not be explicitly illustrated in FIG. **40**.

(549) In at least one embodiment, I/O unit **4006** is configured to transmit and receive communications (e.g., commands, data) from a host processor (not illustrated in FIG. **40**) over system bus **4002**. In at least one embodiment, I/O unit **4006** communicates with host processor directly via system bus **4002** or through one or more intermediate devices such as a memory bridge. In at least one embodiment, I/O unit **4006** may communicate with one or more other processors, such as one or more of PPUs **4000** via system bus **4002**. In at least one embodiment, I/O unit **4006** implements a Peripheral Component Interconnect Express (“PCIe”) interface for communications over a PCIe bus. In at least one embodiment, I/O unit **4006** implements interfaces for communicating with external devices.

(550) In at least one embodiment, I/O unit **4006** decodes packets received via system bus **4002**. In at least one embodiment, at least some packets represent commands configured to cause PPU **4000** to perform various operations. In at least one embodiment, I/O unit **4006** transmits decoded commands to various other units of PPU **4000** as specified by commands. In at least one embodiment, commands are transmitted to front-end unit **4010** and/or transmitted to hub **4016** or other units of PPU **4000** such as one or more copy engines, a video encoder, a video decoder, a power management unit, etc. (not explicitly illustrated in FIG. **40**). In at least one embodiment, I/O unit **4006** is configured to route communications between and among various logical units of PPU **4000**.

(551) In at least one embodiment, a program executed by host processor encodes a command stream in a buffer that provides workloads to PPU **4000** for processing. In at least one embodiment,

a workload comprises instructions and data to be processed by those instructions. In at least one embodiment, buffer is a region in a memory that is accessible (e.g., read/write) by both host processor and PPU **4000**—a host interface unit may be configured to access buffer in a system memory connected to system bus **4002** via memory requests transmitted over system bus **4002** by I/O unit **4006**. In at least one embodiment, host processor writes command stream to buffer and then transmits a pointer to start of command stream to PPU **4000** such that front-end unit **4010** receives pointers to one or more command streams and manages one or more command streams, reading commands from command streams and forwarding commands to various units of PPU **4000**.

(552) In at least one embodiment, front-end unit **4010** is coupled to scheduler unit **4012** that configures various GPCs **4018** to process tasks defined by one or more command streams. In at least one embodiment, scheduler unit **4012** is configured to track state information related to various tasks managed by scheduler unit **4012** where state information may indicate which of GPCs **4018** a task is assigned to, whether task is active or inactive, a priority level associated with task, and so forth. In at least one embodiment, scheduler unit **4012** manages execution of a plurality of tasks on one or more of GPCs **4018**.

(553) In at least one embodiment, scheduler unit **4012** is coupled to work distribution unit **4014** that is configured to dispatch tasks for execution on GPCs **4018**. In at least one embodiment, work distribution unit **4014** tracks a number of scheduled tasks received from scheduler unit **4012** and work distribution unit **4014** manages a pending task pool and an active task pool for each of GPCs **4018**. In at least one embodiment, pending task pool comprises a number of slots (e.g., 32 slots) that contain tasks assigned to be processed by a particular GPC **4018**; active task pool may comprise a number of slots (e.g., 4 slots) for tasks that are actively being processed by GPCs **4018** such that as one of GPCs **4018** completes execution of a task, that task is evicted from active task pool for GPC **4018** and one of other tasks from pending task pool is selected and scheduled for execution on GPC **4018**. In at least one embodiment, if an active task is idle on GPC **4018**, such as while waiting for a data dependency to be resolved, then active task is evicted from GPC **4018** and returned to pending task pool while another task in pending task pool is selected and scheduled for execution on GPC **4018**.

(554) In at least one embodiment, work distribution unit **4014** communicates with one or more GPCs **4018** via XBar **4020**. In at least one embodiment, XBar **4020** is an interconnect network that couples many of units of PPU **4000** to other units of PPU **4000** and can be configured to couple work distribution unit **4014** to a particular GPC **4018**. In at least one embodiment, one or more other units of PPU **4000** may also be connected to XBar **4020** via hub **4016**.

(555) In at least one embodiment, tasks are managed by scheduler unit **4012** and dispatched to one of GPCs **4018** by work distribution unit **4014**. GPC **4018** is configured to process task and generate results. In at least one embodiment, results may be consumed by other tasks within GPC **4018**, routed to a different GPC **4018** via XBar **4020**, or stored in memory **4004**. In at least one embodiment, results can be written to memory **4004** via partition units **4022**, which implement a memory interface for reading and writing data to/from memory **4004**. In at least one embodiment, results can be transmitted to another PPU **4004** or CPU via high-speed GPU interconnect **4008**. In at least one embodiment, PPU **4000** includes, without limitation, a number U of partition units **4022** that is equal to number of separate and distinct memory devices **4004** coupled to PPU **4000**. In at least one embodiment, partition unit **4022** will be described in more detail herein in conjunction with FIG. 42.

(556) In at least one embodiment, a host processor executes a driver kernel that implements an application programming interface (“API”) that enables one or more applications executing on host processor to schedule operations for execution on PPU **4000**. In at least one embodiment, multiple compute applications are simultaneously executed by PPU **4000** and PPU **4000** provides isolation, quality of service (“QoS”), and independent address spaces for multiple compute applications. In at

least one embodiment, an application generates instructions (e.g., in form of API calls) that cause driver kernel to generate one or more tasks for execution by PPU **4000** and driver kernel outputs tasks to one or more streams being processed by PPU **4000**. In at least one embodiment, each task comprises one or more groups of related threads, which may be referred to as a warp. In at least one embodiment, a warp comprises a plurality of related threads (e.g., 32 threads) that can be executed in parallel. In at least one embodiment, cooperating threads can refer to a plurality of threads including instructions to perform task and that exchange data through shared memory. In at least one embodiment, threads and cooperating threads are described in more detail, in accordance with at least one embodiment, in conjunction with FIG. **42**.

(557) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, deep learning application processor is used to train a machine learning model, such as a neural network, to predict or infer information provided to PPU **4000**. In at least one embodiment, deep learning application processor **4000** is used to infer or predict information based on a trained machine learning model (e.g., neural network) that has been trained by another processor or system or by PPU **4000**. In at least one embodiment, PPU **4000** may be used to perform one or more neural network use cases described herein.

(558) In at least one embodiment, PPU **4000** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, PPU **4000** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. PPU **4000** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-39** and **41-43**.

(559) FIG. **41** illustrates a general processing cluster (“GPC”) **4100**, according to at least one embodiment. In at least one embodiment, GPC **4100** is GPC **4018** of FIG. **40**. In at least one embodiment, each GPC **4100** includes, without limitation, a number of hardware units for processing tasks and each GPC **4100** includes, without limitation, a pipeline manager **4102**, a pre-raster operations unit (“PROP”) **4104**, a raster engine **4108**, a work distribution crossbar (“WDX”) **4116**, a memory management unit (“MMU”) **4118**, one or more Data Processing Clusters (“DPCs”) **4106**, and any suitable combination of parts.

(560) In at least one embodiment, operation of GPC **4100** is controlled by pipeline manager **4102**. In at least one embodiment, pipeline manager **4102** manages configuration of one or more DPCs **4106** for processing tasks allocated to GPC **4100**. In at least one embodiment, pipeline manager **4102** configures at least one of one or more DPCs **4106** to implement at least a portion of a graphics rendering pipeline. In at least one embodiment, DPC **4106** is configured to execute a vertex shader program on a programmable streaming multi-processor (“SM”) **4114**. In at least one embodiment, pipeline manager **4102** is configured to route packets received from a work distribution unit to appropriate logical units within GPC **4100**, in at least one embodiment, and some packets may be routed to fixed function hardware units in PROP **4104** and/or raster engine **4108** while other packets may be routed to DPCs **4106** for processing by a primitive engine **4112** or SM **4114**. In at least one embodiment, pipeline manager **4102** configures at least one of DPCs **4106** to implement a neural network model and/or a computing pipeline.

(561) In at least one embodiment, PROP unit **4104** is configured, in at least one embodiment, to route data generated by raster engine **4108** and DPCs **4106** to a Raster Operations (“ROP”) unit in partition unit **4022**, described in more detail above in conjunction with FIG. **40**. In at least one embodiment, PROP unit **4104** is configured to perform optimizations for color blending, organize

pixel data, perform address translations, and more. In at least one embodiment, raster engine **4108** includes, without limitation, a number of fixed function hardware units configured to perform various raster operations, in at least one embodiment, and raster engine **4108** includes, without limitation, a setup engine, a coarse raster engine, a culling engine, a clipping engine, a fine raster engine, a tile coalescing engine, and any suitable combination thereof. In at least one embodiment, setup engine receives transformed vertices and generates plane equations associated with geometric primitive defined by vertices; plane equations are transmitted to coarse raster engine to generate coverage information (e.g., an x, y coverage mask for a tile) for primitive; output of coarse raster engine is transmitted to culling engine where fragments associated with primitive that fail a z-test are culled, and transmitted to a clipping engine where fragments lying outside a viewing frustum are clipped. In at least one embodiment, fragments that survive clipping and culling are passed to fine raster engine to generate attributes for pixel fragments based on plane equations generated by setup engine. In at least one embodiment, output of raster engine **4108** comprises fragments to be processed by any suitable entity such as by a fragment shader implemented within DPC **4106**.

(562) In at least one embodiment, each DPC **4106** included in GPC **4100** comprise, without limitation, an M-Pipe Controller (“MPC”) **4110**; primitive engine **4112**; one or more SMs **4114**; and any suitable combination thereof. In at least one embodiment, MPC **4110** controls operation of DPC **4106**, routing packets received from pipeline manager **4102** to appropriate units in DPC **4106**. In at least one embodiment, packets associated with a vertex are routed to primitive engine **4112**, which is configured to fetch vertex attributes associated with vertex from memory; in contrast, packets associated with a shader program may be transmitted to SM **4114**.

(563) In at least one embodiment, SM **4114** comprises, without limitation, a programmable streaming processor that is configured to process tasks represented by a number of threads. In at least one embodiment, SM **4114** is multi-threaded and configured to execute a plurality of threads (e.g., 32 threads) from a particular group of threads concurrently and implements a Single-Instruction, Multiple-Data (“SIMD”) architecture where each thread in a group of threads (e.g., a warp) is configured to process a different set of data based on same set of instructions. In at least one embodiment, all threads in group of threads execute same instructions. In at least one embodiment, SM **4114** implements a Single-Instruction, Multiple Thread (“SIMT”) architecture wherein each thread in a group of threads is configured to process a different set of data based on same set of instructions, but where individual threads in group of threads are allowed to diverge during execution. In at least one embodiment, a program counter, call stack, and execution state is maintained for each warp, enabling concurrency between warps and serial execution within warps when threads within warp diverge. In another embodiment, a program counter, call stack, and execution state is maintained for each individual thread, enabling equal concurrency between all threads, within and between warps. In at least one embodiment, execution state is maintained for each individual thread and threads executing same instructions may be converged and executed in parallel for better efficiency. At least one embodiment of SM **4114** are described in more detail herein.

(564) In at least one embodiment, MMU **4118** provides an interface between GPC **4100** and memory partition unit (e.g., partition unit **4022** of FIG. 40) and MMU **4118** provides translation of virtual addresses into physical addresses, memory protection, and arbitration of memory requests. In at least one embodiment, MMU **4118** provides one or more translation lookaside buffers (“TLBs”) for performing translation of virtual addresses into physical addresses in memory.

(565) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. 15A and/or 15B. In at least one embodiment, deep learning application processor is used to train a machine learning model, such as a neural network, to predict or infer information provided to GPC **4100**. In at least one embodiment, GPC **4100** is used to infer or predict information based on a trained machine learning

model (e.g., neural network) that has been trained by another processor or system or by GPC **4100**. In at least one embodiment, GPC **4100** may be used to perform one or more neural network use cases described herein.

(566) In at least one embodiment, GPC **4100** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, GPC **4100** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. GPC **4100** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-40** and **42-43**.

(567) FIG. **42** illustrates a memory partition unit **4200** of a parallel processing unit (“PPU”), in accordance with at least one embodiment. In at least one embodiment, memory partition unit **4200** includes, without limitation, a Raster Operations (“ROP”) unit **4202**; a level two (“L2”) cache **4204**; a memory interface **4206**; and any suitable combination thereof memory interface **4206** is coupled to memory. memory interface **4206** may implement 32, 64, 128, 1024-bit data buses, or like, for high-speed data transfer. In at least one embodiment, PPU incorporates U memory interfaces **4206**, one memory interface **4206** per pair of partition units **4200**, where each pair of partition units **4200** is connected to a corresponding memory device. For example, in at least one embodiment, PPU may be connected to up to Y memory devices, such as high bandwidth memory stacks or graphics double-data-rate, version 5, synchronous dynamic random access memory (“GDDR5 SDRAM”).

(568) In at least one embodiment, memory interface **4206** implements a high bandwidth memory second generation (“HBM2”) memory interface and Y equals half U. In at least one embodiment, HBM2 memory stacks are located on same physical package as PPU, providing substantial power and area savings compared with conventional GDDR5 SDRAM systems. In at least one embodiment, each HBM2 stack includes, without limitation, four memory dies and Y equals 4, with each HBM2 stack including two 128-bit channels per die for a total of 8 channels and a data bus width of 1024 bits. In at least one embodiment, memory supports Single-Error Correcting Double-Error Detecting (“SECCDED”) Error Correction Code (“ECC”) to protect data. ECC provides higher reliability for compute applications that are sensitive to data corruption.

(569) In at least one embodiment, PPU implements a multi-level memory hierarchy. In at least one embodiment, memory partition unit **4200** supports a unified memory to provide a single unified virtual address space for central processing unit (“CPU”) and PPU memory, enabling data sharing between virtual memory systems. In at least one embodiment frequency of accesses by a PPU to memory located on other processors is traced to ensure that memory pages are moved to physical memory of PPU that is accessing pages more frequently. In at least one embodiment, high-speed GPU interconnect **4008** supports address translation services allowing PPU to directly access a CPU's page tables and providing full access to CPU memory by PPU.

(570) In at least one embodiment, copy engines transfer data between multiple PPUs or between PPUs and CPUs. In at least one embodiment, copy engines can generate page faults for addresses that are not mapped into page tables and memory partition unit **4200** then services page faults, mapping addresses into page table, after which copy engine performs transfer. In at least one embodiment, memory is pinned (i.e., non-pageable) for multiple copy engine operations between multiple processors, substantially reducing available memory. In at least one embodiment, with hardware page faulting, addresses can be passed to copy engines without regard as to whether memory pages are resident, and copy process is transparent.

(571) Data from memory **4004** of FIG. **40** or other system memory is fetched by memory partition unit **4200** and stored in L2 cache **4204**, which is located on-chip and is shared between various

GPCs, in accordance with at least one embodiment. Each memory partition unit **4200**, in at least one embodiment, includes, without limitation, at least a portion of L2 cache associated with a corresponding memory device. In at least one embodiment, lower level caches are implemented in various units within GPCs. In at least one embodiment, each of SMs **4114** may implement a level one (“L1”) cache wherein L1 cache is private memory that is dedicated to a particular SM **4114** and data from L2 cache **4204** is fetched and stored in each of L1 caches for processing in functional units of SMs **4114**. In at least one embodiment, L2 cache **4204** is coupled to memory interface **4206** and XBar **4020**.

(572) ROP unit **4202** performs graphics raster operations related to pixel color, such as color compression, pixel blending, and more, in at least one embodiment. ROP unit **4202**, in at least one embodiment, implements depth testing in conjunction with raster engine **4108**, receiving a depth for a sample location associated with a pixel fragment from culling engine of raster engine **4108**. In at least one embodiment, depth is tested against a corresponding depth in a depth buffer for a sample location associated with fragment. In at least one embodiment, if fragment passes depth test for sample location, then ROP unit **4202** updates depth buffer and transmits a result of depth test to raster engine **4108**. It will be appreciated that number of partition units **4200** may be different than number of GPCs and, therefore, each ROP unit **4202** can, in at least one embodiment, be coupled to each of GPCs. In at least one embodiment, ROP unit **4202** tracks packets received from different GPCs and determines which that a result generated by ROP unit **4202** is routed to through XBar **4020**.

(573) FIG. **43** illustrates a streaming multi-processor (“SM”) **4300**, according to at least one embodiment. In at least one embodiment, SM **4300** is SM of FIG. **41**. In at least one embodiment, SM **4300** includes, without limitation, an instruction cache **4302**; one or more scheduler units **4304**; a register file **4308**; one or more processing cores (“cores”) **4310**; one or more special function units (“SFUs”) **4312**; one or more load/store units (“LSUs”) **4314**; an interconnect network **4316**; a shared memory/level one (“L1”) cache **4318**; and any suitable combination thereof. In at least one embodiment, a work distribution unit dispatches tasks for execution on general processing clusters (“GPCs”) of parallel processing units (“PPUs”) and each task is allocated to a particular Data Processing Cluster (“DPC”) within a GPC and, if task is associated with a shader program, task is allocated to one of SMs **4300**. In at least one embodiment, scheduler unit **4304** receives tasks from work distribution unit and manages instruction scheduling for one or more thread blocks assigned to SM **4300**. In at least one embodiment, scheduler unit **4304** schedules thread blocks for execution as warps of parallel threads, wherein each thread block is allocated at least one warp. In at least one embodiment, each warp executes threads. In at least one embodiment, scheduler unit **4304** manages a plurality of different thread blocks, allocating warps to different thread blocks and then dispatching instructions from plurality of different cooperative groups to various functional units (e.g., processing cores **4310**, SFUs **4312**, and LSUs **4314**) during each clock cycle.

(574) In at least one embodiment, Cooperative Groups may refer to a programming model for organizing groups of communicating threads that allows developers to express granularity at which threads are communicating, enabling expression of richer, more efficient parallel decompositions. In at least one embodiment, cooperative launch APIs support synchronization amongst thread blocks for execution of parallel algorithms. In at least one embodiment, applications of conventional programming models provide a single, simple construct for synchronizing cooperating threads: a barrier across all threads of a thread block (e.g., `syncthreads( )` function). However, In at least one embodiment, programmers may define groups of threads at smaller than thread block granularities and synchronize within defined groups to enable greater performance, design flexibility, and software reuse in form of collective group-wide function interfaces. In at least one embodiment, Cooperative Groups enables programmers to define groups of threads explicitly at sub-block (i.e., as small as a single thread) and multi-block granularities, and to perform collective operations such as synchronization on threads in a cooperative group.

programming model supports clean composition across software boundaries, so that libraries and utility functions can synchronize safely within their local context without having to make assumptions about convergence. In at least one embodiment, Cooperative Groups primitives enable new patterns of cooperative parallelism, including, without limitation, producer-consumer parallelism, opportunistic parallelism, and global synchronization across an entire grid of thread blocks.

(575) In at least one embodiment, a dispatch unit **4306** is configured to transmit instructions to one or more of functional units and scheduler unit **4304** includes, without limitation, two dispatch units **4306** that enable two different instructions from same warp to be dispatched during each clock cycle. In at least one embodiment, each scheduler unit **4304** includes a single dispatch unit **4306** or additional dispatch units **4306**.

(576) In at least one embodiment, each SM **4300**, in at least one embodiment, includes, without limitation, register file **4308** that provides a set of registers for functional units of SM **4300**. In at least one embodiment, register file **4308** is divided between each of functional units such that each functional unit is allocated a dedicated portion of register file **4308**. In at least one embodiment, register file **4308** is divided between different warps being executed by SM **4300** and register file **4308** provides temporary storage for operands connected to data paths of functional units. In at least one embodiment, each SM **4300** comprises, without limitation, a plurality of L processing cores **4310**. In at least one embodiment, SM **4300** includes, without limitation, a large number (e.g., 128 or more) of distinct processing cores **4310**. In at least one embodiment, each processing core **4310**, in at least one embodiment, includes, without limitation, a fully-pipelined, single-precision, double-precision, and/or mixed precision processing unit that includes, without limitation, a floating point arithmetic logic unit and an integer arithmetic logic unit. In at least one embodiment, floating point arithmetic logic units implement IEEE 754-2008 standard for floating point arithmetic. In at least one embodiment, processing cores **4310** include, without limitation, 64 single-precision (32-bit) floating point cores, 64 integer cores, 32 double-precision (64-bit) floating point cores, and 8 tensor cores.

(577) Tensor cores are configured to perform matrix operations in accordance with at least one embodiment. In at least one embodiment, one or more tensor cores are included in processing cores **4310**. In at least one embodiment, tensor cores are configured to perform deep learning matrix arithmetic, such as convolution operations for neural network training and inferencing. In at least one embodiment, each tensor core operates on a 4×4 matrix and performs a matrix multiply and accumulate operation $D=A \times B+C$, where A, B, C, and D are 4×4 matrices.

(578) In at least one embodiment, matrix multiply inputs A and B are 16-bit floating point matrices and accumulation matrices C and D are 16-bit floating point or 32-bit floating point matrices. In at least one embodiment, tensor cores operate on 16-bit floating point input data with 32-bit floating point accumulation. In at least one embodiment, 16-bit floating point multiply uses 64 operations and results in a full precision product that is then accumulated using 32-bit floating point addition with other intermediate products for a $4 \times 4 \times 4$ matrix multiply. Tensor cores are used to perform much larger two-dimensional or higher dimensional matrix operations, built up from these smaller elements, in at least one embodiment. In at least one embodiment, an API, such as parallel computing platform and application programming interface model 9 C++ API, exposes specialized matrix load, matrix multiply and accumulate, and matrix store operations to efficiently use tensor cores from a parallel computing platform and application programming interface model-C++ program. In at least one embodiment, at parallel computing platform and application programming interface model level, warp-level interface assumes 16×16 size matrices spanning all 32 threads of warp.

(579) In at least one embodiment, each SM **4300** comprises, without limitation, M SFUs **4312** that perform special functions (e.g., attribute evaluation, reciprocal square root, and like). In at least one embodiment, SFUs **4312** include, without limitation, a tree traversal unit configured to traverse a

hierarchical tree data structure. In at least one embodiment, SFUs **4312** include, without limitation, a texture unit configured to perform texture map filtering operations. In at least one embodiment, texture units are configured to load texture maps (e.g., a 2D array of texels) from memory and sample texture maps to produce sampled texture values for use in shader programs executed by SM **4300**. In at least one embodiment, texture maps are stored in shared memory/L1 cache **4318**. In at least one embodiment, texture units implement texture operations such as filtering operations using mip-maps (e.g., texture maps of varying levels of detail), in accordance with at least one embodiment. In at least one embodiment, each SM **4300** includes, without limitation, two texture units.

(580) Each SM **4300** comprises, without limitation, N LSUs **4314** that implement load and store operations between shared memory/L1 cache **4318** and register file **4308**, in at least one embodiment. Each SM **4300** includes, without limitation, interconnect network **4316** that connects each of functional units to register file **4308** and LSU **4314** to register file **4308** and shared memory/L1 cache **4318** in at least one embodiment. In at least one embodiment, interconnect network **4316** is a crossbar that can be configured to connect any of functional units to any of registers in register file **4308** and connect LSUs **4314** to register file **4308** and memory locations in shared memory/L1 cache **4318**.

(581) In at least one embodiment, shared memory/L1 cache **4318** is an array of on-chip memory that allows for data storage and communication between SM **4300** and primitive engine and between threads in SM **4300**, in at least one embodiment. In at least one embodiment, shared memory/L1 cache **4318** comprises, without limitation, 128 KB of storage capacity and is in path from SM **4300** to partition unit. In at least one embodiment, shared memory/L1 cache **4318**, in at least one embodiment, is used to cache reads and writes. In at least one embodiment, one or more of shared memory/L1 cache **4318**, L2 cache, and memory are backing stores.

(582) Combining data cache and shared memory functionality into a single memory block provides improved performance for both types of memory accesses, in at least one embodiment. In at least one embodiment, capacity is used or is usable as a cache by programs that do not use shared memory, such as if shared memory is configured to use half of capacity, texture and load/store operations can use remaining capacity. Integration within shared memory/L1 cache **4318** enables shared memory/L1 cache **4318** to function as a high-throughput conduit for streaming data while simultaneously providing high-bandwidth and low-latency access to frequently reused data, in accordance with at least one embodiment. In at least one embodiment, when configured for general purpose parallel computation, a simpler configuration can be used compared with graphics processing. In at least one embodiment, fixed function graphics processing units are bypassed, creating a much simpler programming model. In general purpose parallel computation configuration, work distribution unit assigns and distributes blocks of threads directly to DPCs, in at least one embodiment. In at least one embodiment, threads in a block execute same program, using a unique thread ID in calculation to ensure each thread generates unique results, using SM **4300** to execute program and perform calculations, shared memory/L1 cache **4318** to communicate between threads, and LSU **4314** to read and write global memory through shared memory/L1 cache **4318** and memory partition unit. In at least one embodiment, when configured for general purpose parallel computation, SM **4300** writes commands that scheduler unit **4304** can use to launch new work on DPCs.

(583) In at least one embodiment, PPU is included in or coupled to a desktop computer, a laptop computer, a tablet computer, servers, supercomputers, a smart-phone (e.g., a wireless, hand-held device), personal digital assistant (“PDA”), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, and more. In at least one embodiment, PPU is embodied on a single semiconductor substrate. In at least one embodiment, PPU is included in a system-on-a-chip (“SoC”) along with one or more other devices such as additional PPUs, memory, a reduced instruction set computer (“RISC”) CPU, a memory management unit (“MMU”), a digital-to-analog

converter (“DAC”), and like.

(584) In at least one embodiment, PPU may be included on a graphics card that includes one or more memory devices. graphics card may be configured to interface with a PCIe slot on a motherboard of a desktop computer. In at least one embodiment, PPU may be an integrated graphics processing unit (“iGPU”) included in chipset of motherboard.

(585) Inference and/or training logic **1515** are used to perform inferencing and/or training operations associated with one or more embodiments. Details regarding inference and/or training logic **1515** are provided herein in conjunction with FIGS. **15A** and/or **15B**. In at least one embodiment, deep learning application processor is used to train a machine learning model, such as a neural network, to predict or infer information provided to SM **4300**. In at least one embodiment, SM **4300** is used to infer or predict information based on a trained machine learning model (e.g., neural network) that has been trained by another processor or system or by SM **4300**. In at least one embodiment, SM **4300** may be used to perform one or more neural network use cases described herein.

(586) In at least one embodiment, SM **4300** or a component thereof executes computer-readable instructions to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores. In at least one embodiment, SM **4300** utilizes computing resources (e.g., CPUs, ASICs, GPUs, FPGAs) to implement inferencing and/or training logic **1515** to perform inferencing and/or training operations associated with one or more embodiments. SM **4300** may be utilized to implement one or more embodiments described elsewhere in this disclosure, such as those described in connection with FIGS. **1-42**.

(587) FIG. **44** illustrates a diagram of unified synchronization, CUDA UMD as writer and reader, according to at least one embodiment. In at least one embodiment, FIG. **44** is implemented in connection with (e.g., in place of, in addition to, as an alternative to) FIGS. **10-11**. In at least one embodiment, an app waiter (e.g., CUDA) calls GetDeviceCapabilities (&w_caps) to obtain a set of device capabilities supported by a UMD waiter. A UMD waiter may determine a set of supported capabilities which can be used to determine a set of constraints on how a cross-UMD sync object is created. In at least one embodiment, UMD waiter specifies what types of synchronization primitives are supported, such as semaphores and syncpoints. In at least one embodiment, UMD waiter sets w_caps to {primitive={semaphore, syncpoint}count={5, 0}} and provides such information to app waiter in any suitable manner. In at least one embodiment, an AttributeList is returned by UMD waiter to app waiter. In at least one embodiment, app waiter sends w_caps to app signaler as AttributeList. In at least one embodiment, signaler sets up a shared fence buffer that is to include fences. In at least one embodiment, app signaler calls GetDeviceCapabilities (&s_caps) to obtain a set of signaler capabilities. In at least one embodiment, UMD signaler fulfills calls to GetDeviceCapabilities (). In at least one embodiment, UMD sets s_caps to indicate a set of supported primitives. In at least one embodiment, s_caps indicates parameters of supported primitives, for example, max count of a semaphore. In at least one embodiment, UMD signaler sets s_caps {primitive={semaphore, syncpoint}count={5, 0}} and returns s_caps as an AttributeList to app signaler. In at least one embodiment, app signaler reconciles AttributeList, finds a suitable primitive for provided combination of signaler and waiter. In at least one embodiment, app signaler calls s_UnifiedSync=UnifiedSyncCreate (attributes) using UnifiedSync which is fulfilled by UnifiedSync. In at least one embodiment, UnifiedSync allocates semaphores if needed (e.g., array of {id, hMem, offset}) and returns s_UnifiedSync to app signaler. In at least one embodiment, app signaler provides UMDImportUnifiedSync (s_UnifiedSync) to UMD signaler. In at least one embodiment, UMD signaler creates a mapping (e.g., GPU MMU mapping), if applicable save away id/index <-> GMMU mapping and associate id/index with UnifiedSyncObject. In at least one embodiment, UMD signaler returns UMD_Signaler object to App signaler. In at least one

embodiment, app signaler setup is performed using techniques described above. In at least one embodiment, app signaler sends UnifiedSync blob using utility APIs to App waiter.

(588) In at least one embodiment, app signaler sends UnifiedSync blob using utility APIs to App waiter. In at least one embodiment, app waiter calls UMDImportUnifiedSync (w_UnifiedSync). In at least one embodiment, UMD waiter creates a mapping (e.g., GMMU mapping) if applicable save away id/index <-> GMMU mapping and associate id/index with UnifiedSyncObject. In at least one embodiment, UMD signaler returns UMD_Waiter object to app signaler. In at least one embodiment, techniques directed towards app waiter setup are described above.

(589) In at least one embodiment, techniques below describe writer/reader synchronized using UMD objects based on UnifiedSync. In at least one embodiment, steps described below can be run in a loop. In at least one embodiment, app signaler, after obtaining UMD_Waiter object, performs tasks on a GPU/SOC engine using UMD signaler. In at least one embodiment, UMD signaler may return status codes indicating success or failure of requested tasks. In at least one embodiment, app signaler calls UMDGenerateUnifiedSyncFence (s_UnifiedSync, UMDObject, UnifiedSyncFence) to generate fences. In at least one embodiment, UMD signaler update fence with semaphore SyncPoint (Index, Value). In at least one embodiment, UMD signaler: finds channel to track from UMDObject; if it is a semaphore, add sema_release (channel, address, value), compose index from address; if it is a syncpoint, add syncpoint release (channel, index, value); and compose <index, value> tuple to return. In at least one embodiment, UMD signaler returns status, fence, a tuple (e.g., {type, index, value}) to app signaler. In at least one embodiment, app signaler sends UnifiedSyncFence using utility APIs provided to app waiter. In at least one embodiment, app waiter calls UMDWaitUnifiedSyncFence (w_UnifiedSync, UMDObject, UnifiedSyncFence). In at least one embodiment, UMD waiter will: if object is semaphore, calculate address for fence received (e.g., address=UMDObject->semaPool+offset); get channel which will wait for fence; if semaphore, add sema_acquire (channel, address, value); if syncpoint, add syncpoint acquire (channel, index, value). In at least one embodiment, UMD waiter returns a status to app waiter. In at least one embodiment, app waiter performs one or more tasks on GPU/SOC engine.

(590) FIG. 45 illustrates a diagram of an intra or inter thread use case, according to at least one embodiment. In at least one embodiment, FIG. 45 illustrates an application 4502; CUDA driver 4504; NvMedia driver 4506; NvSciBuf 4508; and NvRM 4510. In at least one embodiment, application 4502 is a software application that use NvSciBuf (or any other suitable buffer described herein) as part of intra or inter thread operation. In at least one embodiment, memory is shared between at least a CUDA driver 4504 and NvMedia driver 4506.

(591) In at least one embodiment, CUDA constraint settings are implemented in accordance with techniques described in greater detail in this disclosure. In at least one embodiment, application calls an API to create an entry attribute list. In at least one embodiment, application 4502 calls NvSciBuf 4508 using a NvSciBufCreateEmptyAttrlist API and indicates a type (e.g., image) for memory and a reference to a variable for storing an empty attribute list. In at least one embodiment, application 4502 calls NvSciBufCreateEmptyAttrlist (IMAGE, & cuda_attrlist) to create an empty attribute list. In at least one embodiment, application 4502 loops through for all attributes that application 4502 wants to set and sets attribute information. In at least one embodiment, attributes include keys and values. In at least one embodiment, an img_attr variable stores keynames and values for a list of attributes. In at least one embodiment, application 4502 calls a NvSciBufSetAttrs() which sets cuda_attrlist based on image_attrs, which may include also an attr_count that stores how many attributes are in image_attr. In at least one embodiment, application 4502 calls CUDA driver 4504 using an addCUDAConstraints() API that includes at least cuda_attrlist and any other args which CUDA needs. In at least one embodiment, CUDA driver 4504 sets hardware engine information such as a namespace, module identifier, subengine id, version id, and any suitable combination thereof. In at least one embodiment, hardware engine information is stored by CUDA driver 4504 in a hwengine variable or data structure. In at least one

embodiment, CUDA driver **4504** calls `NvSciBuf` **4508** using a `NvSciBufSetEngineList( )` API that includes `cuda_attrlist`, `hwengine`, and an array count (e.g., for one engine, `count=1`). In at least one embodiment, CUDA driver **4504** returns success to application **4502** if no errors.

(592) In at least one embodiment, `NvMedia` constraint settings are described in greater detail here. In at least one embodiment, application **4502** calls `NvMedia` driver **4506** using a `SetNvMediaImageDescriptor( )` API. In at least one embodiment, `NvMedia` driver **4506** calls `NvSciBuf` using a `NvSciBufCreateEmptyAttrlist( )` API to create an empty attribute list identified as `nvmedia_attrlist`. In at least one embodiment, `NvSciBufCreateEmptyAttrlist( )` allows `NvMedia` driver **4506** to specify an object type such as an image or tensor. In at least one embodiment, `NvMedia` driver **4506** iterates through a loop for all attributes that application **4502** wants to set. In at least one embodiment, `image_attr` encodes attribute keys and values for all attributes that application **4502** wants to set. In at least one embodiment, `NvMedia` driver **4506** calls `NvSciBuf` using a `NvSciBufSetAttrs( )` API that sets attributes of `nvmedia_attrlist` using `image_attrs`. In at least one embodiment, `NvMedia` driver **4506** calls `NvSciBuf` using a `NvSciBufCreateDescBlob( )` API which creates a description blob. In at least one embodiment, description blob is `nvm_blob`. In at least one embodiment, UMD info is copied to `nvm_blob`. In at least one embodiment, for all `nvm_engines`, `NvMedia` driver **4506** sets hardware engine information such as namespace, module identifier, subengine identifier, version identifier, and any suitable combination thereof. In at least one embodiment, `NvMedia` driver **4506** sets an engine list using `nvmedia_attrlist` and `hwengine`, including a hardware engine count based on number of `nvm` engines. In at least one embodiment, `NvMedia` driver **4506** returns success to application **4502** if no errors. In at least one embodiment, application **4502** calls `NvSciBuf` using a `NvSciBufSetAttrs( )` API to set general attributes.

(593) In at least one embodiment, after constraints are set, an allocation process is performed. In at least one embodiment, `attrlists` is an array of attribute lists. In at least one embodiment, `attrlists` references an attribute list from CUDA driver **4504** and `NvMedia` driver **4506**. In at least one embodiment, `attrlists` includes one entry that encodes `cuda_attrlist` and a second entry that encodes `nvmedia_attrlist`. In at least one embodiment, application **4502** calls `NvSciBuf` **4508** using a `NvSciBufAppendAttrLists( )` API which includes `attrlists` and returns an unreconciled attribute list `unreconciled_attrlist`. In at least one embodiment, application **4502** calls `NvSciBuf` using a `NvSciBufReconcileAttrList( )` API that takes an unreconciled attribute list to generate a reconciled attribute list. In at least one embodiment, `NvSciBuf` **4508** reconciles attributes using techniques described in connection with FIGS. 1-14. In at least one embodiment, one or more output attributes are computed based on a merged list from reconciled attributes. In at least one embodiment, for an image, output attributes may include pitch, aligned height, and more. In at least one embodiment, `NvSciBuf` loops through for all output attributes to be computed and, in each iteration, gets chip information, reads constraints (e.g., based on chip type, engine list, image type), gets constraints based on image constraints and image pitch alignment, and applies constraints based on merged list image plane width, pitch align value. In at least one embodiment, `NvSciBuf` computes memory allocation attributes from a merged list. In at least one embodiment, reconciled and/or merged attributes are turned to application **4502**. In at least one embodiment, merged attribute is used to call `NvSciBufAllocate( )` using `NvSciBuf`. In at least one embodiment, merged attributes are used to compute memory allocation attributes and memory allocation (e.g., `malloc( )`) APIs are called to obtain a `NvSciBuf` object. In at least one embodiment, memory handle allocation attribute is based on allocation attributes and stored in a memory handle of a `NvSciBuf` object. In at least one embodiment, `NvSciBuf` includes merged attribute list as metadata. In at least one embodiment, a `NvSciBuf` object handle is returned to application **4502**. In at least one embodiment, a success code is returned if no errors (e.g., not enough memory to make allocation) are encountered.

(594) In at least one embodiment, following a success memory allocation by `NvSciBuf` **4508**, application **4502** frees memory allocated to UMD-specific attribute lists such as `cuda_attrlist` and `nvmedia_attrlist` as described in greater detail above.

(595) In at least one embodiment, NvSciBuf memory can be imported by CUDA as a CUDA array. In at least one embodiment, application **4502** calls NvSciBuf using NvSciBufRetrieveMetadata() API by providing NvSciBuf object handle and receiving, in response, and attribute list handle attrlist_h. In at least one embodiment, application **4502** loops through for attributes (e.g., width, height, layout, etc., for CUDA array) and sets key value pairs in an image_attr array. In at least one embodiment, application **4502** calls NvSciBuf using a NvSciBufGetAttrs() API which gets attribute lists based on attrlist_h handle, image_attr, and attr_count (array size of image_attr). In at least one embodiment, application **4502** sets attributes CUDA needs for CUDA array in a cuarraydesc data structure. In at least one embodiment, cuarraydesc includes width, height, offset, etc. properties attributes which are set based on plane values. In at least one embodiment, application **4502** calls CUDA driver **4504** using CreateCUDAArray() API based on cuarraydesc and nvscibufobj_h to obtain a CUDAArray cuArr. In at least one embodiment, CUDA driver **4504** calls BuildCUArrayStructure() API based at least in part on cuarraydesc provided from application **4502**. In at least one embodiment, BuildCUArrayStructure builds CUDAArray. In at least one embodiment, CUDA driver **4504** calls NvSciBuf using NvSciBufGetMemhandle API based on a provided nvscibufobj_h and returns a memhandle which is stored in CUDAArray->rmhandle. In at least one embodiment, CUDA driver **4504** calls MapToGPU based at least in part on CUDAArray->rmhandle and CUDAArray->offset (e.g., set by application **4502**). In at least one embodiment, CUDA driver **4504** calls NvSciBuf using NvSciBufObjDuplicate() to duplicate nvscibufobj_h handle to CUDAArray->nvscibufhandle. In at least one embodiment, CUDAArray is returned to application **4502**.

(596) In at least one embodiment NvSciBuf can be imported as a NvMediaImage. In at least one embodiment, application **4502** calls NvMedia driver **4506** using CreateNvMediaImage() API based at least in part on a NvSciBuf buffer handle, nvscibufobj_h. In at least one embodiment, NvMedia driver **4506** gets a cookie count. In at least one embodiment, NvMedia driver **4506** calls NvSciBuf using a NvSciBufGetCookeCount() API based at least in part on nvscibufobj_h. In at least one embodiment, NvMedia driver **4506** specifies in NvSciBufGetCookeCount() API a value that indicates it as a UMD and gets a cookie count in return. In at least one embodiment, NvMedia driver **4506** calls NvSciBuf **4508** using NvSciBufGetCookies() API based at least in part on nvscibufobj_h and cookie count to obtain a nvmbjlist. If nvmbjlist is non-empty, NvMedia driver **4506** calls NvSciBuf **4508** using NvSciBufRetrieveMetdata using nvscibufobj_h to obtain metadata, which may be returned as an attribute list handle, attrlist_h. In at least one embodiment, NvMedia driver **4506** calls NvSciBuf **4508** using NvSciBufGetDescBlobCount() API using attrlist_h and indicating UMD NVMEDIA to get a blob count. In at least one embodiment, multiple blobs of same type may exist. In at least one embodiment, NvMedia driver **4506** calls NvSciBuf GetDescBlobs() API based at least in part on attrlist_h, UMD NVMEDIA, and blob count to obtain NVMDescblobs. In at least one embodiment, NvMedia driver **4506** calls NvSciBuf **4508** using NvSciBufGetAttrs() based at least in part on attrlist_h and attribute count to obtain image_attrs. In at least one embodiment, NvMedia driver **4506** builds a NvMedia image using NVMDescblobs and attributes. In at least one embodiment, NvMedia driver **4506** calls NvSciBuf **4508** using NvSciBufObjDuplicate() API to duplicate nvscibufobj_h to produce a copy nvm_nvscibuf_h. In at least one embodiment, NvMImage->nvscibuf_h is set to nvm_nvscibuf_h. In at least one embodiment, NvMImage->rmhandle stores a memhandle that is obtained from NvSciBuf **4508** using NvSciBufGetMemhandle() API based at least in part on nvscibufobj_h. In at least one embodiment, NvMedia driver **4506** calls NvSciBuf **4508** using NvSciBufSetCookie() based at least in part on nvscibufobj_h, UMD NVMEDIA, NvMImage. In at least one embodiment, nvmediaimage is set based on NvMImage. In at least one embodiment, NvMedia driver **4506** returns nvmediaimage to application **4502**. In at least one embodiment, application **4502** calls NvMedia driver **4506** using CreateDeviceMappings() based on nvm_image. In at least one embodiment, NvMedia driver **4506** calls NvRM **4510** using DoDevicePinnings() API based at

least in part on nvmmimage.

(597) In at least one embodiment, as part of a runtime, application **4502** submits a CUDA kernel CUDAArray to CUDA driver **4504**. In at least one embodiment, application **4502** submits a nvmm_image to IO engine. Sequences of run times are merely illustrative and non-limiting.

(598) In at least one embodiment, teardown sequence can occur in any suitable order. In at least one embodiment, a CUDA array CUDAArray is detached by application **4502** calling CUDA driver **4504** using a free() API on CUDAArray. In at least one embodiment, CUDA driver **4504** unmaps to GPU CUDAArray->NvRmMemHandle, which is submitted to NvRM. In at least one embodiment, CUDA driver **4504** calls NvSciBuf **4508** using NvSciBufFree on CUDAArray->nvscibuf_h) which is processed by NvSciBuf **4508**. In at least one embodiment, CUDA driver **4504** calls free() on CUDAArray after unmapping nvRmMemHandle and calling NvSciBufFree on nvscibuf_h. In at least one embodiment, CUDA driver **4504** returns a success or error code to application **4502**.

(599) In at least one embodiment, a NvMedia image is detached. In at least one embodiment, application **4502** calls NvMedia driver **4506** using UnmaptoDevices() based at least in part on nvmm_image. In at least one embodiment, application **4502** calls NvMedia driver **4506** using free() API on nvmm_image. In at least one embodiment, NvMedia driver **4506** calls NvSciBuf **4508** using NvSciBufRemoveCookie() API based on nvmm_image and nvmm_image->nvscibuf_h. In at least one embodiment, NvMedia driver **4506** calls NvSciBuf **4508** calls NvSciBufFree() based at least in part on nvmm_image->nvscibuf_h. In at least one embodiment, NvMedia driver **4506** calls free() on nvmm_image. In at least one embodiment, NvMedia driver **4506** returns a success or error code to application **4502**.

(600) In at least one embodiment, application **4502** calls NvSciBuf **4508** to free NvSciBufObj based at least in part on nvscibuf_h. In at least one embodiment, NvSciBuf **4508** obtains a handle and gets NvSciBufObj from corresponding handle nvscibuf_h. If no references to NvSciBufObj exist (e.g., refcount is equal to 0), NvSciBuf **4508** may call free() on NvSciBufObj->attrlist and NvRmMemhandleFree() on NvSciBufObj->NvRmMemHandle. In at least one embodiment, once all memory objects within NvSciBufObj are freed, NvSciBuf **4508** calls free() on NvSciBufObj.

(601) At least one embodiment of the disclosure can be described in view of the following clauses:

(602) Clause 1. A processor, comprising: one or more circuits to allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores.

(603) Clause 2. The processor of clause 1, wherein the at least two heterogeneous processing cores comprises a central processing unit (CPU) and a graphics processing unit (GPU).

(604) Clause 3. The processor of any of clauses 1-2, wherein the one or more attributes indicates whether to use system memory or video memory.

(605) Clause 4. The processor of any of clauses 1-3, wherein the video memory is accessible by a discrete graphics processing unit (dGPU).

(606) Clause 5. The processor of any of clauses 1-4, the one or more circuits to allocate the memory to the at least two heterogeneous processing cores are to process the one or more attributes to determine a set of constraints on how the memory is allocated.

(607) Clause 6. The processor of any of clauses 1-5, wherein the memory is allocated in a manner that to be interpreted as a first data object by a first heterogeneous processing core of the at least two heterogeneous processing cores and to be interpreted as a second data object by a second heterogeneous processing core of the at least two heterogeneous processing cores.

(608) Clause 7. The processor of any of clauses 1-6, wherein the one or more circuits are to further: obtain different one or more attributes associated with how the at least two heterogeneous processing cores support coordinating access to the memory; determine a manner in which to initialize a synchronization object to coordinate access to the memory based at least in part on the

different one or more attributes; and provide the at least two heterogeneous processing cores access to the synchronization object.

(609) Clause 8. The processor of any of clauses 1-7, wherein the synchronization object is a semaphore.

(610) Clause 9. A system, comprising one or more memories to store instructions that, as a result of execution by one or more processors, cause the system to: allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores.

(611) Clause 10. The system of clause 9, wherein the at least two heterogeneous processing cores comprises at least a portion of the one or more processors.

(612) Clause 11. The system of any of clauses 9-10, wherein the instructions to cause the system to allocate the memory to the at least two heterogeneous processing cores are instructions that, as a result of execution by the one or more processors, cause the system to process the one or more attributes to determine a manner in which to allocate the memory.

(613) Clause 12. The system of any of clauses 9-11, wherein the manner in which to allocate the memory satisfies constraints imposed by attributes provided by the at least two heterogeneous processing cores.

(614) Clause 13. The system of any of clauses 9-12, wherein the memory maps to a parallel computing platform and application programming interface model object.

(615) Clause 14. The system of any of clauses 9-13, wherein the instructions to allocate the memory are instructions that, as a result of execution by the one or more processors, cause the system to provide access to the memory via a handle that is to be interpreted by the at least two heterogeneous processing cores.

(616) Clause 15. The system of any of clauses 9-14, wherein the handle is interpreted as a first data object by a first heterogeneous processing core of the at least two heterogeneous processing cores and interpreted as a second data object by a second heterogeneous processing core of the at least two heterogeneous processing cores.

(617) Clause 16. The system of any of clauses 9-15, wherein the one or more memories are to store instructions that, as a result of execution by the one or more processors, cause the system to: obtain different one or more attributes associated with how the at least two heterogeneous processing cores support coordinating access to the memory; determine a manner in which to initialize a signal to coordinate access to the memory based at least in part on the different one or more attributes; and provide the at least two heterogeneous processing cores access to the signal.

(618) Clause 17. The system of any of clauses 9-16, wherein the different one or more attributes encodes types of synchronization primitives supported by the at least two heterogeneous processing cores.

(619) Clause 18. A method, comprising: allocating memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with the API based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores.

(620) Clause 19. The method of clause 18, wherein the at least two heterogeneous processing cores comprises a first central processing unit (CPU) and second CPU of different instruction set architectures.

(621) Clause 20. The method of any of clauses 18-19, wherein the first CPU supports an ARM instruction set architecture.

(622) Clause 21. The method of any of clauses 18-20, wherein the second CPU supports an x86 instruction set architecture.

(623) Clause 22. The method of any of clauses 18-21, wherein allocating the memory to the at least two heterogeneous processing cores comprises: determining, based at least in part on the one or

more attributes, a set of allocation semantics associated with the at least two heterogeneous processing cores; and determining a manner in which to allocate the memory that satisfies one or more constraints imposed by the set of allocation semantics.

(624) Clause 23. The method of any of clauses 18-22, wherein the memory is interpreted as a tensor by a first core of the at least two heterogeneous processing cores and is interpreted as a texture by a second core of the at least two heterogeneous processing cores.

(625) Clause 24. The method of any of clauses 18-23, wherein the one or more attributes correspond to the at least two heterogeneous processing cores.

(626) Clause 25. The method of any of clauses 18-24, wherein the memory is exposed, by the API, as a handle to be interpreted by the at least two heterogeneous processing cores.

(627) Clause 26. The method of any of clauses 18-25, further comprising: obtaining different one or more attributes associated with how the at least two heterogeneous processing cores support coordinating access to the memory; determining a manner in which to initialize a signal to coordinate access to the memory based at least in part on the different one or more attributes; and providing the at least two heterogeneous processing cores access to the signal.

(628) Clause 27. The method of any of clauses 18-26, wherein providing the at least two heterogeneous processing cores access to the signal comprises providing a handle to the signal with signaling and waiting semantics to be interpreted by the at least two heterogeneous processing cores.

(629) Clause 28. A machine-readable medium having stored thereon an application programming interface (API), which if performed by one or more processors, cause the one or more processors to at least: allocate memory to at least two heterogeneous processing cores in response to performing one or more instructions associated with the API based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores.

(630) Clause 29. The machine-readable medium of clause 28, wherein the at least two heterogeneous processing cores comprises an accelerator.

(631) Clause 30. The machine-readable medium of any of clauses 28-29, wherein the accelerator is a programmable vision accelerator (PVA).

(632) Clause 31. The machine-readable medium of any of clauses 28-30, wherein the machine-readable medium comprises instructions which, if performed by the one or more processors, cause the one or more processors to store data to the memory as a first type of data object and read the data from the memory as a second type of data object.

(633) Clause 32. The machine-readable medium of any of clauses 28-31, wherein the first type of data object is an image and the second type of data object is a tensor.

(634) Clause 33. The machine-readable medium of any of clauses 28-32, wherein the API to allocate the memory, if performed by the one or more processors, causes the one or more processors to provide a first handle to the memory and a second handle to the one or more attributes.

(635) Clause 34. A processor, comprising: one or more circuits to create a signal to be used to coordinate at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores.

(636) Clause 35. The processor of clause 34, wherein the signal is to be used to coordinate execution of computer-readable instructions between the at least two heterogeneous processing cores.

(637) Clause 36. The processor of any of clauses 34-35, wherein the signal is to be used to coordinate access to memory between the at least two heterogeneous processing cores.

(638) Clause 37. The processor of any of clauses 34-36, wherein the signal is to be interpreted as a first synchronization primitive by a first heterogeneous processing core of the at least two heterogeneous processing cores and to be interpreted as a second synchronization primitive by a

second heterogeneous processing core of the at least two heterogeneous processing cores.

(639) Clause 38. The processor of any of clauses 34-37, wherein the first synchronization primitive is a semaphore and the second synchronization primitive is a fence.

(640) Clause 39. The processor of any of clauses 34-38, wherein the at least two heterogeneous processing cores comprises a central processing unit and a graphics processing unit.

(641) Clause 40. The processor of any of clauses 34-39, wherein the one or more circuits are to further: allocate memory to be shared between the at least two heterogeneous processing cores support coordinating access to the memory; and coordinate access to the memory using the signal.

(642) Clause 41. The processor of any of clauses 34-40, wherein the one or more circuits are to coordinate access to the memory using the signal by at least causing a first heterogeneous processing cores to wait on a second heterogeneous processing cores.

(643) Clause 42. A system, comprising one or more memories to store instructions that, as a result of execution by one or more processors, cause the system to: create a signal to be used to coordinate at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores.

(644) Clause 43. The system of clause 42, wherein the signal is to be used to synchronize execution of the at least two heterogeneous processing cores.

(645) Clause 44. The system of any of clauses 42-43, wherein the signal is to be used to synchronize data access between the at least two heterogeneous processing cores.

(646) Clause 45. The system of any of clauses 42-44, wherein the instructions to cause the system to create a signal to be used to coordinate at least two heterogeneous processing cores are instructions that, as a result of execution by the one or more processors, cause the system to process the one or more attributes to determine a manner in which to create the signal.

(647) Clause 46. The system of any of clauses 42-45, wherein the manner in which to create the signal satisfies constraints imposed by attributes of the at least two heterogeneous processing cores through the one or more APIs.

(648) Clause 47. The system of any of clauses 42-46, wherein the instructions to create the signal are instructions that, as a result of execution by the one or more processors, cause the system to provide access to the signal via a handle that is to be interpreted by the at least two heterogeneous processing cores.

(649) Clause 48. The system of any of clauses 42-47, wherein the handle is interpreted as a first synchronization object by a first heterogeneous processing core of the at least two heterogeneous processing cores and interpreted as a second synchronization object by a second heterogeneous processing core of the at least two heterogeneous processing cores.

(650) Clause 49. The system of any of clauses 42-48, wherein the one or more memories are to store instructions that, as a result of execution by the one or more processors, cause the system to: obtain different one or more attributes associated the at least two heterogeneous processing cores; determine a set of constrains on memory allocation based at least in part on the different one or more attributes; and allocate memory to be shared by the at least two heterogeneous processing cores, according to the set of constraints.

(651) Clause 50. The system of any of clauses 42-49, wherein the memory is to be interpreted as a first data object by a first heterogeneous processing core of the at least two heterogeneous processing cores and to be interpreted as a second object by a second heterogeneous processing core of the at least two heterogeneous processing cores.

(652) Clause 51. A method, comprising: creating a signal to be used to coordinate at least two heterogeneous processing cores in response to performing one or more instructions associated with one or more application programming interfaces (APIs) based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores.

(653) Clause 52. The method of clause 51, wherein the signal is to be used to coordinate scheduling

of executable code between the at least two heterogeneous processing cores.

(654) Clause 53. The method of any of clauses 51-52, wherein the signal is to be used to coordinate access to memory between the at least two heterogeneous processing cores.

(655) Clause 54. The method of any of clauses 51-53, wherein the signal is implemented to be interpreted as a first synchronization primitive by a first heterogeneous processing core of the at least two heterogeneous processing cores and to be interpreted as a second synchronization primitive by a second heterogeneous processing core of the at least two heterogeneous processing cores.

(656) Clause 55. The method of any of clauses 51-54, wherein the first synchronization primitive is a semaphore and the second synchronization primitive is a syncpoint.

(657) Clause 56. The method of any of clauses 51-55, wherein the at least two heterogeneous processing cores comprises a central processing unit and a graphics processing unit.

(658) Clause 57. The method of any of clauses 51-56, wherein the one or more circuits are to further: allocate memory to be shared between the at least two heterogeneous processing cores support coordinating access to the memory; and coordinate access to the memory using the signal.

(659) Clause 58. The method of any of clauses 51-57, wherein the one or more circuits are to coordinate access to the memory using the signal by at least causing a first heterogeneous processing cores to wait on a second heterogeneous processing cores.

(660) Clause 59. A machine-readable medium having stored thereon one or more application programming interfaces (APIs), which if performed by one or more processors, cause the one or more processors to at least: create a signal to be used to coordinate at least two heterogeneous processing cores in response to performing one or more instructions associated with the one or more APIs based, at least in part, on one or more attributes associated with the at least two heterogeneous processing cores.

(661) Clause 60. The machine-readable medium of clause 59, wherein the signal is to be used to coordinate execution of computer-readable instructions between the at least two heterogeneous processing cores.

(662) Clause 61. The machine-readable medium of any of clauses 59-60, wherein the signal is to be used by a first heterogeneous processing cores of the at least two heterogeneous processing cores to block access to memory accessible to a second heterogeneous processing cores of the at least two heterogeneous processing cores.

(663) Clause 62. The machine-readable medium of any of clauses 59-61, wherein the signal is to be interpreted as a first synchronization primitive by a first heterogeneous processing core of the at least two heterogeneous processing cores and to be interpreted as a second synchronization primitive by a second heterogeneous processing core of the at least two heterogeneous processing cores.

(664) Clause 63. The machine-readable medium of any of clauses 59-62, wherein the signal is to be used by a first heterogeneous processing core to signal a second first heterogeneous processing core waiting on the signal.

(665) Clause 64. The machine-readable medium of any of clauses 59-63, wherein the one or more circuits are to further: allocate memory to be shared between the at least two heterogeneous processing cores support coordinating access to the memory; and coordinate access to the memory using the signal.

(666) Clause 65. The machine-readable medium of any of clauses 59-64, wherein memory is to store one or more images and the signal is to coordinate access to the memory between a camera and a graphics processing unit.

(667) In at least one embodiment, a signal is implemented as any suitable synchronization primitive or synchronization mechanism. In at least one embodiment, a signal is used to coordinate access to data (e.g., memory) and/or execution of code between two or more heterogeneous processing cores. In at least one embodiment, non-limiting examples of a signal includes: semaphores; syncpoints;

fences; shared and exclusive locks; events; mutexes; spinlocks; critical sections; and more.

(668) In at least one embodiment, a single semiconductor platform may refer to a sole unitary semiconductor-based integrated circuit or chip. In at least one embodiment, multi-chip modules may be used with increased connectivity which simulate on-chip operation, and make substantial improvements over utilizing a conventional central processing unit (“CPU”) and bus implementation. In at least one embodiment, various modules may also be situated separately or in various combinations of semiconductor platforms per desires of user.

(669) In at least one embodiment, computer programs in form of machine-readable executable code or computer control logic algorithms are stored in main memory **2104** and/or secondary storage. Computer programs, if executed by one or more processors, enable system **2100** to perform various functions in accordance with at least one embodiment. memory **2104**, storage, and/or any other storage are possible examples of computer-readable media. In at least one embodiment, secondary storage may refer to any suitable storage device or system such as a hard disk drive and/or a removable storage drive, representing a floppy disk drive, a magnetic tape drive, a compact disk drive, digital versatile disk (“DVD”) drive, recording device, universal serial bus (“USB”) flash memory, etc. In at least one embodiment, architecture and/or functionality of various previous figures are implemented in context of CPU **2102**; parallel processing system **2112**; an integrated circuit capable of at least a portion of capabilities of both CPU **2102**; parallel processing system **2112**; a chipset (e.g., a group of integrated circuits designed to work and sold as a unit for performing related functions, etc.); and any suitable combination of integrated circuit(s).

(670) In at least one embodiment, architecture and/or functionality of various previous figures are implemented in context of a general computer system, a circuit board system, a game console system dedicated for entertainment purposes, an application-specific system, and more. In at least one embodiment, computer system **2100** may take form of a desktop computer, a laptop computer, a tablet computer, servers, supercomputers, a smart-phone (e.g., a wireless, hand-held device), personal digital assistant (“PDA”), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, a mobile phone device, a television, workstation, game consoles, embedded system, and/or any other type of logic.

(671) In at least one embodiment, parallel processing system **2112** includes, without limitation, a plurality of parallel processing units (“PPUs”) **2114** and associated memories **2116**. In at least one embodiment, PPUs **2114** are connected to a host processor or other peripheral devices via an interconnect **2118** and a switch **2120** or multiplexer. In at least one embodiment, parallel processing system **2112** distributes computational tasks across PPUs **2114** which can be parallelizable—for example, as part of distribution of computational tasks across multiple graphics processing unit (“GPU”) thread blocks. In at least one embodiment, memory is shared and accessible (e.g., for read and/or write access) across some or all of PPUs **2114**, although such shared memory may incur performance penalties relative to use of local memory and registers resident to a PPU **2114**. In at least one embodiment, operation of PPUs **2114** is synchronized through use of a command such as `_syncthreads( )` wherein all threads in a block (e.g., executed across multiple PPUs **2114**) to reach a certain point of execution of code before proceeding.

(672) Other variations are within spirit of present disclosure. Thus, while disclosed techniques are susceptible to various modifications and alternative constructions, certain illustrated embodiments thereof are shown in drawings and have been described above in detail. It should be understood, however, that there is no intention to limit disclosure to specific form or forms disclosed, but on contrary, intention is to cover all modifications, alternative constructions, and equivalents falling within spirit and scope of disclosure, as defined in appended claims.

(673) Use of terms “a” and “an” and “the” and similar referents in context of describing disclosed embodiments (especially in context of following claims) are to be construed to cover both singular and plural, unless otherwise indicated herein or clearly contradicted by context, and not as a definition of a term. Terms “comprising,” “having,” “including,” and “containing” are to be

construed as open-ended terms (meaning “including, but not limited to,”) unless otherwise noted. term “connected,” when unmodified and referring to physical connections, is to be construed as partly or wholly contained within, attached to, or joined together, even if there is something intervening. Recitation of ranges of values herein are merely intended to serve as a shorthand method of referring individually to each separate value falling within range, unless otherwise indicated herein and each separate value is incorporated into specification as if it were individually recited herein. use of term “set” (e.g., “a set of items”) or “subset” unless otherwise noted or contradicted by context, is to be construed as a nonempty collection comprising one or more members. Further, unless otherwise noted or contradicted by context, term “subset” of a corresponding set does not necessarily denote a proper subset of corresponding set, but subset and corresponding set may be equal.

(674) Conjunctive language, such as phrases of form “at least one of A, B, and C,” or “at least one of A, B and C,” unless specifically stated otherwise or otherwise clearly contradicted by context, is otherwise understood with context as used in general to present that an item, term, etc., may be either A or B or C, or any nonempty subset of set of A and B and C. For instance, in illustrative example of a set having three members, conjunctive phrases “at least one of A, B, and C” and “at least one of A, B and C” refer to any of following sets: {A}, {B}, {C}, {A, B}, {A, C}, {B, C}, {A, B, C}. Thus, such conjunctive language is not generally intended to imply that certain embodiments require at least one of A, at least one of B and at least one of C each to be present. In addition, unless otherwise noted or contradicted by context, term “plurality” indicates a state of being plural (e.g., “a plurality of items” indicates multiple items). number of items in a plurality is at least two, but can be more when so indicated either explicitly or by context. Further, unless stated otherwise or otherwise clear from context, phrase “based on” means “based at least in part on” and not “based solely on.”

(675) Operations of processes described herein can be performed in any suitable order unless otherwise indicated herein or otherwise clearly contradicted by context. In at least one embodiment, a process such as those processes described herein (or variations and/or combinations thereof) is performed under control of one or more computer systems configured with executable instructions and is implemented as code (e.g., executable instructions, one or more computer programs or one or more applications) executing collectively on one or more processors, by hardware or combinations thereof. In at least one embodiment, code is stored on a computer-readable storage medium, for example, in form of a computer program comprising a plurality of instructions executable by one or more processors. In at least one embodiment, a computer-readable storage medium is a non-transitory computer-readable storage medium that excludes transitory signals (e.g., a propagating transient electric or electromagnetic transmission) but includes non-transitory data storage circuitry (e.g., buffers, cache, and queues) within transceivers of transitory signals. In at least one embodiment, code (e.g., executable code or source code) is stored on a set of one or more non-transitory computer-readable storage media having stored thereon executable instructions (or other memory to store executable instructions) that, when executed (i.e., as a result of being executed) by one or more processors of a computer system, cause computer system to perform operations described herein. set of non-transitory computer-readable storage media, in at least one embodiment, comprises multiple non-transitory computer-readable storage media and one or more of individual non-transitory storage media of multiple non-transitory computer-readable storage media lack all of code while multiple non-transitory computer-readable storage media collectively store all of code. In at least one embodiment, executable instructions are executed such that different instructions are executed by different processors—for example, a non-transitory computer-readable storage medium store instructions and a main central processing unit (“CPU”) executes some of instructions while a graphics processing unit (“GPU”) executes other instructions. In at least one embodiment, different components of a computer system have separate processors and different processors execute

different subsets of instructions.

(676) Accordingly, in at least one embodiment, computer systems are configured to implement one or more services that singly or collectively perform operations of processes described herein and such computer systems are configured with applicable hardware and/or software that enable performance of operations. Further, a computer system that implements at least one embodiment of present disclosure is a single device and, in another embodiment, is a distributed computer system comprising multiple devices that operate differently such that distributed computer system performs operations described herein and such that a single device does not perform all operations.

(677) Use of any and all examples, or exemplary language (e.g., “such as”) provided herein, is intended merely to better illuminate embodiments of disclosure and does not pose a limitation on scope of disclosure unless otherwise claimed. No language in specification should be construed as indicating any non-claimed element as essential to practice of disclosure.

(678) All references, including publications, patent applications, and patents, cited herein are hereby incorporated by reference to same extent as if each reference were individually and specifically indicated to be incorporated by reference and were set forth in its entirety herein.

(679) In description and claims, terms “coupled” and “connected,” along with their derivatives, may be used. It should be understood that these terms may be not intended as synonyms for each other. Rather, in particular examples, “connected” or “coupled” may be used to indicate that two or more elements are in direct or indirect physical or electrical contact with each other. “Coupled” may also mean that two or more elements are not in direct contact with each other, but yet still co-operate or interact with each other.

(680) Unless specifically stated otherwise, it may be appreciated that throughout specification terms such as “processing,” “computing,” “calculating,” “determining,” or like, refer to action and/or processes of a computer or computing system, or similar electronic computing device, that manipulate and/or transform data represented as physical, such as electronic, quantities within computing system's registers and/or memories into other data similarly represented as physical quantities within computing system's memories, registers or other such information storage, transmission or display devices.

(681) In a similar manner, term “processor” may refer to any device or portion of a device that processes electronic data from registers and/or memory and transform that electronic data into other electronic data that may be stored in registers and/or memory. As non-limiting examples, “processor” may be a CPU or a GPU. A “computing platform” may comprise one or more processors. As used herein, “software” processes may include, for example, software and/or hardware entities that perform work over time, such as tasks, threads, and intelligent agents. Also, each process may refer to multiple processes, for carrying out instructions in sequence or in parallel, continuously or intermittently. terms “system” and “method” are used herein interchangeably insofar as system may embody one or more methods and methods may be considered a system.

(682) In present document, references may be made to obtaining, acquiring, receiving, or inputting analog or digital data into a subsystem, computer system, or computer-implemented machine. process of obtaining, acquiring, receiving, or inputting analog and digital data can be accomplished in a variety of ways such as by receiving data as a parameter of a function call or a call to an application programming interface. In some implementations, process of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a serial or parallel interface. In another implementation, process of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a computer network from providing entity to acquiring entity. References may also be made to providing, outputting, transmitting, sending, or presenting analog or digital data. In various examples, process of providing, outputting, transmitting, sending, or presenting analog or digital data can be accomplished by transferring data as an input or output parameter of a function call, a parameter of

an application programming interface or interprocess communication mechanism.
(683) Although discussion above sets forth example implementations of described techniques, other architectures may be used to implement described functionality, and are intended to be within scope of this disclosure. Furthermore, although specific distributions of responsibilities are defined above for purposes of discussion, various functions and responsibilities might be distributed and divided in different ways, depending on circumstances.

(684) Furthermore, although subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that subject matter claimed in appended claims is not necessarily limited to specific features or acts described. Rather, specific features and acts are disclosed as exemplary forms of implementing the claims.

Claims

1. A processor, comprising: one or more circuits to perform one or more application programming interfaces (APIs) to cause different types of memory to be allocated to at least two heterogeneous processing cores based, at least in part, on one or more different attributes associated with the at least two heterogeneous processing cores.
2. The processor of claim 1, wherein the at least two heterogeneous processing cores comprise a central processing unit and a graphics processing unit.
3. The processor of claim 1, wherein the one or more different attributes indicates whether to use system memory or video memory.
4. The processor of claim 3, wherein the video memory is accessible by a discrete graphics processing unit.
5. The processor of claim 1, wherein the one or more circuits to allocate the memory to the at least two heterogeneous processing cores are to process the one or more different attributes to determine a set of constraints on how the different types of memory are to be allocated.
6. The processor of claim 5, wherein the different types of memory are to be allocated in a manner that to be interpreted as a first data object by a first heterogeneous processing core of the at least two heterogeneous processing cores and to be interpreted as a second data object by a second heterogeneous processing core of the at least two heterogeneous processing cores.
7. The processor of claim 1, wherein the one or more circuits are to further: obtain the one or more different attributes associated with how the at least two heterogeneous processing cores support coordinating access to the different types of memory; determine a manner in which to initialize a synchronization object to coordinate access to the different types of memory based at least in part on the one or more different attributes; and provide the at least two heterogeneous processing cores access to the synchronization object.
8. The processor of claim 7, wherein the synchronization object is a semaphore.
9. A system, comprising one or more memories to store instructions that, as a result of execution by one or more processors, cause the system to: perform one or more application programming interfaces (APIs) to cause different types of memory to be allocated to at least two heterogeneous processing cores based, at least in part, on one or more different attributes associated with the at least two heterogeneous processing cores.
10. The system of claim 9, wherein the at least two heterogeneous processing cores comprise at least a portion of the one or more processors.
11. The system of claim 9, wherein the instructions to cause the system to allocate the different types of memory to the at least two heterogeneous processing cores are instructions that, as a result of execution by the one or more processors, cause the system to process the one or more different attributes to determine a manner in which to allocate the different types of memory.
12. The system of claim 11, wherein the manner in which to allocate the different types of memory satisfies constraints imposed by the one or more different attributes of the at least two

heterogeneous processing cores through the API.

13. The system of claim 9, wherein the different types of memory map to a parallel computing platform and application programming interface model object.

14. The system of claim 9, wherein the instructions to allocate the different types of memory are instructions that, as a result of execution by the one or more processors, cause the system to provide access to the different types of memory via a handle that is to be interpreted by the at least two heterogeneous processing cores.

15. The system of claim 14, wherein the handle is interpreted as a first data object by a first heterogeneous processing core of the at least two heterogeneous processing cores and interpreted as a second data object by a second heterogeneous processing core of the at least two heterogeneous processing cores.

16. The system of claim 9, wherein the one or more memories are to store instructions that, as a result of execution by the one or more processors, cause the system to: obtain the one or more different attributes associated with how the at least two heterogeneous processing cores support coordinating access to the different types of memory; determine a manner in which to initialize a signal to coordinate access to the different types of memory based at least in part on the one or more different attributes; and provide the at least two heterogeneous processing cores access to the signal.

17. The system of claim 16, wherein the one or more different attributes encodes types of synchronization primitives supported by the at least two heterogeneous processing cores.

18. A method, comprising: performing one or more application programming interfaces (APIs) to cause different types of memory to be allocated to at least two heterogeneous processing cores based, at least in part, on one or more different attributes associated with the at least two heterogeneous processing cores.

19. The method of claim 18, wherein the at least two heterogeneous processing cores comprise a first central processing unit (CPU) and second CPU of different instruction set architectures.

20. The method of claim 19, wherein the first CPU supports an ARM instruction set architecture.

21. The method of claim 20, wherein the second CPU supports an x86 instruction set architecture.

22. The method of claim 18, wherein allocating the memory to the at least two heterogeneous processing cores comprises: determining, based at least in part on the one or more different attributes, a set of allocation semantics associated with the at least two heterogeneous processing cores; and determining a manner in which to allocate the different types of memory that satisfy one or more constraints imposed by the set of allocation semantics.

23. The method of claim 22, wherein the memory is interpreted as a tensor by a first core of the at least two heterogeneous processing cores and is interpreted as a texture by a second core of the at least two heterogeneous processing cores.

24. The method of claim 18, wherein the one or more different attributes correspond to the at least two heterogeneous processing cores.

25. The method of claim 18, wherein the different types of memory are exposed, by the API, as a handle to be interpreted by the at least two heterogeneous processing cores.

26. The method of claim 18, further comprising: obtaining the one or more different attributes associated with how the at least two heterogeneous processing cores support coordinating access to the different types of memory; determining a manner in which to initialize a signal to coordinate access to the different types of memory based at least in part on the one or more different attributes; and providing the at least two heterogeneous processing cores access to the signal.

27. The method of claim 26, wherein providing the at least two heterogeneous processing cores access to the signal comprises providing a handle to the signal with signaling and waiting semantics to be interpreted by the at least two heterogeneous processing cores.

28. A non-transitory machine-readable medium having stored thereon an application programming interface (API), which if performed by one or more processors, cause the one or more processors to

at least: allocate different types of memory to at least two heterogeneous processing cores based, at least in part, on one or more different attributes associated with the at least two heterogeneous processing cores.

29. The non-transitory machine-readable medium of claim 28, wherein the at least two heterogeneous processing cores comprise an accelerator.

30. The non-transitory machine-readable medium of claim 29, wherein the accelerator is a programmable vision accelerator.

31. The non-transitory machine-readable medium of claim 28, wherein the machine-readable medium comprises instructions which, if performed by the one or more processors, cause the one or more processors to store data to the different types of memory as a first type of data object and read the data from the different types of memory as a second type of data object.

32. The non-transitory machine-readable medium of claim 31, wherein the first type of data object is an image and the second type of data object is a tensor.

33. The non-transitory machine-readable medium of claim 28, wherein the API, if performed by the one or more processors, causes the one or more processors to provide a first handle to the different types of memory and a second handle to the one or more different attributes.

34. A processor, comprising: one or more circuits to perform one or more application programming interfaces (APIs) to create a signal to be used to coordinate allocating different types of memory to at least two heterogeneous processing cores based, at least in part, on one or more different attributes associated with the at least two heterogeneous processing cores.

35. The processor of claim 34, wherein the signal is to be used to coordinate execution of computer-readable instructions between the at least two heterogeneous processing cores.

36. The processor of claim 34, wherein the signal is to be used to coordinate access to the different types of memory between the at least two heterogeneous processing cores.

37. The processor of claim 34, wherein the signal is to be interpreted as a first synchronization primitive by a first heterogeneous processing core of the at least two heterogeneous processing cores and to be interpreted as a second synchronization primitive by a second heterogeneous processing core of the at least two heterogeneous processing cores.

38. The processor of claim 37, wherein the first synchronization primitive is a semaphore and the second synchronization primitive is a fence.

39. The processor of claim 34, wherein the at least two heterogeneous processing cores comprise a central processing unit and a graphics processing unit.

40. The processor of claim 34, wherein the one or more circuits are to further: allocate memory to be shared between the at least two heterogeneous processing cores support coordinating access to the memory; and coordinate access to the memory using the signal.

41. The processor of claim 40, wherein the one or more circuits are to coordinate access to the memory using the signal by at least causing a first heterogeneous processing cores to wait on a second heterogeneous processing cores.

42. A system, comprising one or more memories to store instructions that, as a result of execution by one or more processors, cause the system to: perform one or more application programming interfaces (APIs) to create a signal to be used to coordinate allocating different types of memory to at least two heterogeneous processing cores based, at least in part, on one or more different attributes associated with the at least two heterogeneous processing cores.

43. The system of claim 42, wherein the signal is to be used to synchronize execution of the at least two heterogeneous processing cores.

44. The system of claim 42, wherein the signal is to be used to synchronize data access between the at least two heterogeneous processing cores.

45. The system of claim 42, wherein the instructions to cause the system to create a signal to be used to coordinate at least two heterogeneous processing cores are instructions that, as a result of execution by the one or more processors, cause the system to process the one or more different

attributes to determine a manner in which to create the signal.

46. The system of claim 45, wherein the manner in which to create the signal satisfies constraints imposed by attributes of the at least two heterogeneous processing cores through the API.

47. The system of claim 42, wherein the instructions to create the signal are instructions that, as a result of execution by the one or more processors, cause the system to provide access to the signal via a handle that is to be interpreted by the at least two heterogeneous processing cores.

48. The system of claim 47, wherein the handle is interpreted as a first synchronization object by a first heterogeneous processing core of the at least two heterogeneous processing cores and interpreted as a second synchronization object by a second heterogeneous processing core of the at least two heterogeneous processing cores.

49. The system of claim 42, wherein the one or more memories are to store instructions that, as a result of execution by the one or more processors, cause the system to: obtain the one or more different attributes associated with the at least two heterogeneous processing cores; determine a set of constraints on memory allocation based at least in part on the one or more different attributes; and allocate memory to be shared by the at least two heterogeneous processing cores, according to the set of constraints.

50. The system of claim 49, wherein the memory is to be interpreted as a first data object by a first heterogeneous processing core of the at least two heterogeneous processing cores and to be interpreted as a second object by a second heterogeneous processing core of the at least two heterogeneous processing cores.

51. A method, comprising: performing one or more application programming interfaces (APIs) to create a signal to be used to coordinate allocating different types of memory to at least two heterogeneous processing cores based, at least in part, on one or more different attributes associated with the at least two heterogeneous processing cores.

52. The method of claim 51, wherein the signal is to be used to coordinate scheduling of executable code between the at least two heterogeneous processing cores.

53. The method of claim 51, wherein the signal is to be used to coordinate access to the different types of memory between the at least two heterogeneous processing cores.

54. The method of claim 51, wherein the signal is implemented to be interpreted as a first synchronization primitive by a first heterogeneous processing core of the at least two heterogeneous processing cores and to be interpreted as a second synchronization primitive by a second heterogeneous processing core of the at least two heterogeneous processing cores.

55. The method of claim 54, wherein the first synchronization primitive is a semaphore and the second synchronization primitive is a syncpoint.

56. The method of claim 51, wherein the at least two heterogeneous processing cores comprise a central processing unit and a graphics processing unit.

57. The method of claim 51, wherein one or more circuits are to further: allocate memory to be shared between the at least two heterogeneous processing cores support coordinating access to the memory; and coordinate access to the memory using the signal.

58. The method of claim 57, wherein the one or more circuits are to coordinate access to the memory using the signal by at least causing a first heterogeneous processing cores to wait on a second heterogeneous processing cores.

59. A non-transitory machine-readable medium having stored thereon one or more application programming interfaces (APIs), which if performed by one or more processors, cause the one or more processors to at least: create a signal to be used to coordinate allocating different types of memory to at least two heterogeneous processing cores based, at least in part, on one or more different attributes associated with the at least two heterogeneous processing cores.

60. The non-transitory machine-readable medium of claim 59, wherein the signal is to be used to coordinate execution of computer-readable instructions between the at least two heterogeneous processing cores.

61. The non-transitory machine-readable medium of claim 59, wherein the signal is to be used by a first heterogeneous processing cores of the at least two heterogeneous processing cores to block access to memory accessible to a second heterogeneous processing cores of the at least two heterogeneous processing cores.
62. The non-transitory machine-readable medium of claim 59, wherein the signal is to be interpreted as a first synchronization primitive by a first heterogeneous processing core of the at least two heterogeneous processing cores and to be interpreted as a second synchronization primitive by a second heterogeneous processing core of the at least two heterogeneous processing cores.
63. The non-transitory machine-readable medium of claim 59, wherein the signal is to be used by a first heterogeneous processing core to signal a second first heterogeneous processing core waiting on the signal.
64. The non-transitory machine-readable medium of claim 59, wherein the one or more processors are to further: allocate memory to be shared between the at least two heterogeneous processing cores support coordinating access to the memory; and coordinate access to the memory using the signal.
65. The non-transitory machine-readable medium of claim 64, wherein memory is to store one or more images and the signal is to coordinate access to the memory between a camera and a graphics processing unit.
-