

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent

12405857

Kind Code

B2

Date of Patent

September 02, 2025

Inventor(s)

Gole; Abhijeet Prakash et al.

Methods and systems for raid protection in zoned solid-state drives

Abstract

Methods and systems for a storage environment are provided. One method includes splitting storage of a plurality of zoned solid-state drives (ZNS SSDs) into a plurality of physical zones (PZones) across a plurality of independent media units of each ZNS SSD, the PZones visible to a first tier RAID (redundant array of independent disks) layer; generating a plurality of RAID zones (RZones), each RZone having a plurality of PZones; presenting one or more RZones to a second tier RAID layer by the first tier RAID layer for processing read and write requests using the plurality of ZNS SSDs; and utilizing, by the first tier RAID layer, a parity PZone at each ZNS SSD for storing parity information corresponding to data written in one or more PZone corresponding to a RZone presented to the second tier RAID layer and storing the parity information in a single parity ZNS SSD.

Inventors: Gole; Abhijeet Prakash (Cupertino, CA), Sen; Sourav (Cupertino, CA), Smith; Mark (Cupertino, CA), Ting; Daniel Wang-Woei (Palo Alto, CA)

Applicant: NETAPP, INC. (San Jose, CA)

Family ID: 80446042

Assignee: NETAPP, INC. (San Jose, CA)

Appl. No.: 18/209613

Filed: June 14, 2023

Prior Publication Data

Document Identifier

Publication Date

US 20230325279 A1

Oct. 12, 2023

Related U.S. Application Data

continuation parent-doc US 17727511 20220422 US 11698836 child-doc US 18209613

Publication Classification

Int. Cl.: **G06F11/10** (20060101); **G06F3/06** (20060101); **G06F12/10** (20160101)

U.S. Cl.:

CPC **G06F11/1076** (20130101); **G06F3/0619** (20130101); **G06F3/0644** (20130101);
G06F3/0689 (20130101); **G06F12/10** (20130101); G06F2212/657 (20130101)

Field of Classification Search

CPC: G06F (3/0689); G06F (11/108); G06F (11/1076); G06F (2212/7201); G06F (12/10);
G06F (11/2094); G06F (11/1088); G06F (3/0619); G06F (3/0644); G06F (2212/657);
G11C (29/76)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5664187	12/1996	Burkes et al.	N/A	N/A
6101615	12/1999	Lyons	N/A	N/A
8074021	12/2010	Miller et al.	N/A	N/A
8463991	12/2012	Colgrove et al.	N/A	N/A
8549222	12/2012	Kleiman et al.	N/A	N/A
8775868	12/2013	Colgrove et al.	N/A	N/A
8832373	12/2013	Colgrove et al.	N/A	N/A
8850108	12/2013	Hayes et al.	N/A	N/A
8862820	12/2013	Colgrove et al.	N/A	N/A
9003144	12/2014	Hayes et al.	N/A	N/A
9021297	12/2014	Hayes et al.	N/A	N/A
9134917	12/2014	Kimmel et al.	N/A	N/A
9201600	12/2014	Hayes et al.	N/A	N/A
9218244	12/2014	Hayes et al.	N/A	N/A
9229808	12/2015	Colgrove et al.	N/A	N/A
9244769	12/2015	Colgrove et al.	N/A	N/A
9367243	12/2015	Hayes et al.	N/A	N/A
9483346	12/2015	Davis et al.	N/A	N/A
9495255	12/2015	Davis et al.	N/A	N/A
9525738	12/2015	Hayes et al.	N/A	N/A
9563506	12/2016	Hayes et al.	N/A	N/A
9588842	12/2016	Sanvido et al.	N/A	N/A
9594633	12/2016	Colgrove et al.	N/A	N/A
9672125	12/2016	Botes et al.	N/A	N/A
9672905	12/2016	Gold et al.	N/A	N/A
9798477	12/2016	Botes et al.	N/A	N/A
9880899	12/2017	Davis et al.	N/A	N/A
9934089	12/2017	Hayes et al.	N/A	N/A

9967342	12/2017	Colgrove et al.	N/A	N/A
10095701	12/2017	Faibish et al.	N/A	N/A
10180879	12/2018	Colgrove et al.	N/A	N/A
10248516	12/2018	Sanvido et al.	N/A	N/A
10303547	12/2018	Hayes et al.	N/A	N/A
10353777	12/2018	Bernat et al.	N/A	N/A
10372506	12/2018	Baptist et al.	N/A	N/A
10379763	12/2018	Colgrove et al.	N/A	N/A
10387247	12/2018	Baptist et al.	N/A	N/A
10387250	12/2018	Resch et al.	N/A	N/A
10387256	12/2018	Dhuse et al.	N/A	N/A
10402266	12/2018	Kirkpatrick et al.	N/A	N/A
10417092	12/2018	Brennan et al.	N/A	N/A
10432233	12/2018	Colgrove et al.	N/A	N/A
10437673	12/2018	Baptist et al.	N/A	N/A
10437678	12/2018	Resch	N/A	N/A
10452289	12/2018	Colgrove et al.	N/A	N/A
10467107	12/2018	Abrol et al.	N/A	N/A
10489256	12/2018	Hayes et al.	N/A	N/A
10503598	12/2018	Trichardt et al.	N/A	N/A
10521120	12/2018	Miller et al.	N/A	N/A
10530862	12/2019	Isely et al.	N/A	N/A
10534661	12/2019	Resch	N/A	N/A
10572176	12/2019	Davis et al.	N/A	N/A
10579450	12/2019	Khadiwala et al.	N/A	N/A
10606700	12/2019	Alnafoosi et al.	N/A	N/A
10613974	12/2019	Dreier et al.	N/A	N/A
10656871	12/2019	Peake	N/A	N/A
10657000	12/2019	Resch	N/A	N/A
10671480	12/2019	Hayes et al.	N/A	N/A
RE48222	12/2019	Colgrove et al.	N/A	N/A
10776200	12/2019	McVay et al.	N/A	N/A
10776204	12/2019	Resch et al.	N/A	N/A
10810083	12/2019	Colgrove et al.	N/A	N/A
10817375	12/2019	Colgrove et al.	N/A	N/A
10838834	12/2019	Sanvido et al.	N/A	N/A
10860424	12/2019	Dhuse et al.	N/A	N/A
10891192	12/2020	Brennan et al.	N/A	N/A
RE48448	12/2020	Colgrove et al.	N/A	N/A
11269778	12/2021	Kanteti	N/A	N/A
11340987	12/2021	Gole	N/A	G06F 3/0644
11442646	12/2021	Agarwal	N/A	N/A
11698836	12/2022	Gole	714/764	G06F 11/108
11709734	12/2022	Kaynak et al.	N/A	N/A
11789611	12/2022	Singh et al.	N/A	N/A
11797377	12/2022	Gole	N/A	N/A
11803329	12/2022	Doucette et al.	N/A	N/A
11816359	12/2022	Gole et al.	N/A	N/A

11861195	12/2023	Gorobets et al.	N/A	N/A
11861231	12/2023	Gole et al.	N/A	N/A
12242337	12/2024	Gole	N/A	N/A
2006/0129873	12/2005	Hafner et al.	N/A	N/A
2006/0242539	12/2005	Kang et al.	N/A	N/A
2010/0332401	12/2009	Prahlad et al.	N/A	N/A
2012/0084506	12/2011	Colgrove et al.	N/A	N/A
2012/0151118	12/2011	Flynn et al.	N/A	N/A
2012/0198152	12/2011	Terry	711/E12.001	G06F 11/1092
2013/0151646	12/2012	Chidambaram et al.	N/A	N/A
2014/0281227	12/2013	Herron et al.	N/A	N/A
2015/0169244	12/2014	Asnaashari et al.	N/A	N/A
2015/0199151	12/2014	Klemm et al.	N/A	N/A
2016/0239397	12/2015	Thomas	N/A	G06F 3/0689
2016/0313943	12/2015	Hashimoto et al.	N/A	N/A
2016/0342470	12/2015	Cudak et al.	N/A	N/A
2017/0124345	12/2016	Christiansen et al.	N/A	N/A
2017/0220264	12/2016	Sokolov et al.	N/A	N/A
2019/0004964	12/2018	Kanno	N/A	N/A
2019/0018788	12/2018	Yoshida et al.	N/A	N/A
2019/0087111	12/2018	Lingarajappa	N/A	G06F 3/0631
2019/0114096	12/2018	Trika	N/A	G06F 3/0619
2019/0278663	12/2018	Mehta et al.	N/A	N/A
2020/0089407	12/2019	Baca et al.	N/A	N/A
2020/0394112	12/2019	Gupta et al.	N/A	N/A
2020/0409589	12/2019	Bennett et al.	N/A	N/A
2020/0409601	12/2019	Helmick et al.	N/A	N/A
2021/0026717	12/2020	Krasner et al.	N/A	N/A
2021/0081273	12/2020	Helmick	N/A	G06F 9/30189
2021/0081330	12/2020	Bennett et al.	N/A	N/A
2021/0132827	12/2020	Helmick et al.	N/A	N/A
2021/0132860	12/2020	Kou et al.	N/A	N/A
2021/0303188	12/2020	Bazarsky et al.	N/A	N/A
2021/0326048	12/2020	Karr	N/A	N/A
2021/0334006	12/2020	Singh et al.	N/A	N/A
2021/0334041	12/2020	Bennett	N/A	G06F 3/0652
2021/0397517	12/2020	Sharma	N/A	G06F 3/0631
2022/0027051	12/2021	Kant et al.	N/A	N/A
2022/0137844	12/2021	Goss et al.	N/A	N/A
2022/0197553	12/2021	Benhanokh et al.	N/A	N/A
2022/0227051	12/2021	Regehly	N/A	N/A
2022/0229596	12/2021	Jung	N/A	N/A

2022/0244869	12/2021	Kanteti	N/A	N/A
2022/0291838	12/2021	Gorobets et al.	N/A	N/A
2023/0022758	12/2022	Sears et al.	N/A	N/A
2023/0082636	12/2022	Zhu et al.	N/A	N/A
2023/0333740	12/2022	Singh et al.	N/A	N/A
2023/0350858	12/2022	Karr et al.	N/A	N/A
2024/0028262	12/2023	Gole et al.	N/A	N/A
2024/0045760	12/2023	Gole	N/A	N/A
2025/0060911	12/2024	Gole et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
1343087	12/2002	EP	N/A

OTHER PUBLICATIONS

A. A. McEwan and I. Mir, "An Embedded FTL for SSD RAID," 2015 Euromicro Conference on Digital System Design, Madeira, Portugal, 2015, pp. 575-582, doi: 10.1109/DSD.2015.39.

keywords: {Ash;Reliability;Arrays;Metadata;Hardware;Random access memory;SSD storage;embedded FTL;SSD RAID;wear leveling;FPGA}, (Year: 2015). cited by examiner
Dholakia A., et al., "A New Intra-disk Redundancy Scheme for High-Reliability RAID Storage Systems in the Presence of Unrecoverable Errors," ACM Transactions on Storage, May 2008, vol. 4 (1), Article 1,42 pages. cited by applicant

European Search Report for Application No. EP22157793 mailed on Jul. 19, 2022, 16 pages. cited by applicant

International Preliminary Report on Patentability for Application No. PCT/US2021/028879, mailed on Oct. 25, 2022, 6 pages. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2021/028879, mailed on Aug. 5, 2021, 8 pages. cited by applicant

International Search Report and Written Opinion for Patent Application No. PCT/US2022/049431, mailed on Mar. 3, 2023, 13 pages. cited by applicant

Mao B., et al., "HPDA: A Hybrid Parity-Based Disk Array for Enhanced Performance and reliability," ACM Transactions on Storage (TOS), vol. 8(1), Publication [online], Feb. 2012 [retrieved Apr. 4, 2016]. Retrieved from the Internet: URL:

<http://or.nsfc.gov.cn/bitstream/00001903-5/90177/1/1000003549834.pdf> , 20 pages. cited by applicant

Mao B., et al., "HPDA: A Hybrid Parity-Based Disk Array for Enhanced Performance and Reliability," Retrieved from Internet URL: <https://www.researchgate.net/publication/224140602>, May 2020; 13 pages. cited by applicant

NetApp, Inc., "Data ONTAP7.3 Data Protection Online Backup and Recovery Guide," Feb. 22, 2011, Part No. 210-05212_A0, 432 pages. cited by applicant

NetApp, Inc., "Data ONTAP@7.3 Active/Active Configuration Guide," Jun. 16, 2011, Part No. 210-05247_A0, 214 pages. cited by applicant

NetApp, Inc., "Data ONTAP7.3 Archive and Compliance Management Guide," Mar. 4, 2010, Part No. 210-04827_A0, 180 pages. cited by applicant

NetApp, Inc., "Data ONTAP7.3 Block Access Management Guide for iSCSI and FC," Mar. 4, 2010, Part No. 210-04752_B0, 202 pages. cited by applicant

NetApp, Inc., "Data ONTAP7.3 Data Protection Tape Backup and Recovery Guide," Jan. 15, 2010, Part No. 210-04762_A0, 142 pages. cited by applicant

NetApp, Inc., "Data ONTAP7.3 Documentation Roadmap," Jul. 9, 2008, Part No. 210-04229_A0,

8 pages. cited by applicant

NetApp, Inc., "Data ONTAP7.3 File Access and Protocols Management Guide," Sep. 10, 2009, Part No. 210-04505_B0, 382 pages. cited by applicant

NetApp, Inc., "Data ONTAP7.3 MultiStore Management Guide," Mar. 4, 2010, Part No. 210-04855_A0, 144 pages. cited by applicant

NetApp, Inc., "Data ONTAP7.3 Network Management Guide," Jan. 15, 2010, Part No. 210-04757_A0, 222 pages. cited by applicant

NetApp, Inc., "Data ONTAP7.3 Software Setup Guide," Nov. 4, 2010, Part No. 210-05045_A0, 116 pages. cited by applicant

NetApp, Inc., "Data ONTAP7.3 Storage Efficiency Management Guide," Mar. 4, 2010, Part No. 210-04856_A0, 76 pages. cited by applicant

NetApp, Inc., "Data ONTAP7.3 Storage Management Guide," May 3, 2012, Part No. 210-04766_B0, 356 pages. cited by applicant

NetApp, Inc., "Data ONTAP7.3 System Administration Guide," Nov. 11, 2010, Part No. 210-05043_A0, 350 pages. cited by applicant

NetApp, Inc., "Data ONTAP7.3 Upgrade Guide," Nov. 11, 2010, Part No. 210-05042_A0, 200 pages. cited by applicant

NetApp, Inc., "Data ONTAP7.3.7 Release Notes," May 31, 2012, Part No. 215-06916_A0, 182 pages. cited by applicant

NetApp, Inc., "Data ONTAP7.3 Core Commands Quick Reference," Jun. 2008, Part No. 215-03893_A0, 1 page. cited by applicant

NetApp, Inc., "Notices," 2010, Part No. 215-05705_A0, 46 pages. cited by applicant

NetApp, Inc., "V-Series Systems Hardware Maintenance Guide," Jul. 2006, Part No. 210-00975_A0, 202 pages. cited by applicant

NetApp, Inc., "V-Series Systems Implementation Guide for Hitachi Storage," Dec. 2009, Part No. 210-04694_A0, 66 pages. cited by applicant

NetApp, Inc., "V-Series Systems Installation Requirements and Reference Guide," Oct. 2010, Part No. 210-05064_A0, 214 pages. cited by applicant

NetApp, Inc., "V-Series Systems MetroCluster Guide," Jul. 2009, Part No. 210-04515_A0, 80 pages. cited by applicant

Non-Final Office Action mailed on Feb. 16, 2023 for U.S. Appl. No. 17/650,940, filed Feb. 14, 2022, 24 pages. cited by applicant

Non-Final Office Action mailed on Apr. 18, 2023 for U.S. Appl. No. 17/456,012, filed Nov. 22, 2021, 34 pages. cited by applicant

Non-Final Office Action mailed on Mar. 30, 2023 for U.S. Appl. No. 17/494,684, filed Oct. 5, 2021, 8 pages. cited by applicant

Non-Final Office Action mailed on May 15, 2023 for U.S. Appl. No. 17/650,936, filed Feb. 14, 2022, 19 pages. cited by applicant

Notice of Allowance mailed on Apr. 14, 2023 for U.S. Appl. No. 17/727,511, filed Apr. 22, 2022, 2 pages. cited by applicant

Notice of Allowance mailed on Jun. 1, 2023 for U.S. Appl. No. 17/727,511, filed Apr. 22, 2022, 2 pages. cited by applicant

Notice of Allowance mailed on Jun. 16, 2023 for U.S. Appl. No. 16/858,019, filed Apr. 24, 2020, 10 pages. cited by applicant

Notice of Allowance mailed on Mar. 1, 2023 for U.S. Appl. No. 17/727,511, filed Apr. 22, 2022, 15 pages. cited by applicant

Notice of Allowance mailed on May 24, 2023 for U.S. Appl. No. 17/727,511, filed Apr. 22, 2022, 2 pages. cited by applicant

Notice of Allowance mailed on Dec. 20, 2022 for U.S. Appl. No. 16/858,019, filed Apr. 24, 2020, 11 pages. cited by applicant

Notice of Allowance mailed on Dec. 14, 2022 for U.S. Appl. No. 17/727,511, filed Apr. 22, 2022 , 12 pages. cited by applicant
NVM Express Base Specification, Mar. 9, 2020, Revision 1.4a; NVM Express Workgroup, 405 pages. cited by applicant
Notice of Allowance mailed on Aug. 30, 2023 for U.S. Appl. No. 17/456,012, filed Nov. 22, 2021, 10 pages. cited by applicant
U.S. Appl. No. 17/727,511, filed Apr. 22, 2022. cited by applicant
U.S. Appl. No. 17/192,606, filed Mar. 4, 2021. cited by applicant
Non-Final Office Action mailed on May 22, 2024 for U.S. Appl. No. 18/377,441, filed Oct. 6, 2023, 10 pages. cited by applicant
Non-Final Office Action mailed on Sep. 19, 2024 for U.S. Appl. No. 18/343,149, filed Jun. 28, 2023, 16 pages. cited by applicant
Notice of Allowance mailed on Sep. 13, 2024 for U.S. Appl. No. 18/377,441, filed Oct. 6, 2023, 05 pages. cited by applicant
Notice of Allowance mailed on Jul. 19, 2023 for U.S. Appl. No. 17/650,940, filed Feb. 14, 2022, 9 pages. cited by applicant
Notice of Allowance mailed on Jun. 26, 2023 for U.S. Appl. No. 17/494,684, filed Oct. 5, 2021, 8 pages. cited by applicant
Notice of Allowance mailed on Jul. 5, 2024 for U.S. Appl. No. 18/474,988, filed Sep. 26, 2023, 02 pages. cited by applicant
Notice of Allowance mailed on Jun. 26, 2024 for U.S. Appl. No. 18/474,988, filed Sep. 26, 2023, 08 pages. cited by applicant
Corrected Notice of Allowance mailed on Feb. 4, 2025 for U.S. Appl. No. 18/343,149, filed Jun. 28, 2023, 15 pages. cited by applicant
Notice of Allowance mailed on Feb. 21, 2025 for U.S. Appl. No. 18/343,149, filed Jun. 28, 2023, 11 pages. cited by applicant
Notice of Allowance mailed on Jan. 31, 2025 for U.S. Appl. No. 18/377,441, filed Oct. 6, 2023, 02 pages. cited by applicant
Notice of Allowance mailed on Mar. 12, 2025 for U.S. Appl. No. 18/343,149, filed Jun. 28, 2023, 11 pages. cited by applicant

Primary Examiner: McMahon; Daniel F.

Attorney, Agent or Firm: LOZA & LOZA, LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION (1) This patent application claims priority of and is a continuation of co-pending U.S. patent application Ser. No. 17/727,511, filed on Apr. 22, 2022, which claims priority of and is a continuation of U.S. patent application Ser. No. 17/192,606, filed on Mar. 4, 2021, now U.S. Pat. No. 11,340,987, issued on May 24, 2022, the disclosures of which are incorporated herein by reference in their entirety.

TECHNICAL FIELD

(1) The present disclosure relates to storage environments and more particularly, for providing RAID (redundant array of independent (or inexpensive) disks) protection in zoned solid-state drives.

BACKGROUND

(2) Various forms of storage systems are used today. These forms include direct attached storage

- (DAS) network attached storage (NAS) systems, storage area networks (SANs), and others. Network storage systems are commonly used for a variety of purposes, such as providing multiple users with access to shared data, backing up data and others.
- (3) A storage system typically includes at least one computing system executing a storage operating system for storing and retrieving data on behalf of one or more client computing systems (“clients”). The storage operating system stores and manages shared data containers in a set of mass storage devices operating in a group of a storage sub-system. The storage devices (may also be referred to as “disks”) within a storage system are typically organized as one or more groups (or arrays), wherein each group is operated as a RAID.
- (4) Most RAID implementations enhance reliability/integrity of data storage through redundant writing of data “stripes” across a given number of physical drives in the RAID group and storing parity data associated with striped data in dedicated parity drives. A storage device may fail in a storage sub-system. Data can be lost when one or more storage devices fail. The parity data is used to protect against loss of data in a RAID group.
- (5) RAID6 and RAID-DP (RAID-Dual Parity) type protection is typically employed to protect RAID groups against dual drive failures. Both RAID6 and RAID-DP employ two redundant storage drives to store dual parity data. Double failure protection by dual parity RAID includes ability to continue providing data after two drives have failed or a single drive has failed and one of the other drives in the RAID group encounters an uncorrectable read error.
- (6) Conventional dual parity RAID schemes allocate at least two dedicated storage drives for storing parity data. This additional cost of dual parity protection is undesirable, especially when the storage drives are high-capacity SSDs and the RAID group contains fewer drives. For example, using 2 out of 8 drives of a RAID group to store parity data significantly reduces the overall storage capacity and increases the cost of storing parity data. Continuous efforts are being made to develop technology for providing dual parity data protection (e.g., RAID 6 and RAID-DP type protection) without having to use more than one parity drive or a RAID group.
-

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The foregoing features and other features will now be described with reference to the drawings of the various aspects. In the drawings, the same components have the same reference numerals. The illustrated aspects are intended to illustrate, but not to limit the present disclosure. The drawings include the following Figures:
- (2) FIG. 1A shows an example of an operating environment for the various aspects disclosed herein;
- (3) FIG. 1B shows illustrates configuration of ZNS (Zone Namespace) SSDs (solid state drives), according to one aspect of the present disclosure;
- (4) FIG. 1C provides another example of the ZNS SSD configuration, according to one aspect of the present disclosure;
- (5) FIG. 1D shows an example architecture for using ZNS SSDs, according to one aspect of the present disclosure;
- (6) FIG. 1E shows an example of a “commit” operation for a physical zone (PZone) of a ZNS SSD;
- (7) FIG. 1F shows an example of a “commit” operation for a RAID zone (RZone) of a ZNS SSD;
- (8) FIG. 1G shows a process for initializing PZones and RZones of a ZNS SSD, according to one aspect of the present disclosure;
- (9) FIG. 2A shows a process for parity generation, according to one aspect of the present disclosure;
- (10) FIG. 2B illustrates a parity generation example, according to one aspect of the present

disclosure;

(11) FIG. 2C shows a process for parity overwrite, according to one aspect of the present disclosure;

(12) FIG. 2D illustrates a parity overwrite example, according to one aspect of the present disclosure;

(13) FIG. 2E shows a process for executing a read operation, according to one aspect of the present disclosure;

(14) FIG. 2F shows a process for data reconstruction, according to one aspect of the present disclosure;

(15) FIG. 3 shows an example of a storage operating system, used according to one aspect of the present disclosure; and

(16) FIG. 4 shows an example of a processing system, used according to one aspect of the present disclosure.

DETAILED DESCRIPTION

(17) In one aspect, innovative technology is provided to enable data protection against dual failures using parity information (also referred to as parity data) that is stored in one parity drive and in a plurality of data drives within a RAID (Redundant Array of Independent (or Inexpensive) Disks) group (or array, used interchangeably throughout this specification). Unlike conventional solutions provided by RAID-6 and RAID-DP, dual redundant parity drives are not used or needed for certain type of failure conditions. The disclosed technical solution saves cost because additional parity drives are not used, and the available storage capacity of a RAID group increases because two drives are not used to just store parity data.

(18) In one aspect, the technology disclosed herein uses zoned namespace solid state drives (“ZNS SSDs”). A ZNS SSD has individual media units (“Mus”) that operate independent of each other to store information. Storage space at each ZNS SSD is exposed as zones. The zones are configured using the independent MUs, which enables the MUs to operate as individual drives of a RAID group. A first tier RAID layer configures the storage space of ZNS SSDs into physical zones (“PZones”) spread uniformly across the MUs. Each MU is configured to include a plurality of PZones. The first tier RAID layer configures a plurality of RAID zones (“RZones”), each having a plurality of PZones. The RZones are presented to other layers, e.g., a tier 2 RAID layer that interfaces with a file system to process read and write requests. The tier 2 RAID layer and the file system manager only see the RZone and the tier 1 layer manages data at the PZone level.

(19) Parity is determined by XORing data stored across a horizontal stripe having a plurality of PZones. The parity data is stored at a single parity ZNS SSD and also within a parity PZone of each ZNS SSD. If a block or a MU fails, then the parity data stored at the individual ZNS SSD or the parity drive is used to reconstruct data. This provides RAID-6 and RAID-DP type parity protection without having to use two or more dedicated parity drives. Details regarding the innovative technology of the present disclosure are provided below.

(20) As a preliminary note, the terms “component”, “module”, “system,” and the like as used herein are intended to refer to a computer-related entity, either software-executing general-purpose processor, hardware, firmware and a combination thereof. For example, a component may be, but is not limited to being, a process running on a hardware processor, a hardware processor, an object, an executable, a thread of execution, a program, and/or a computer.

(21) By way of illustration, both an application running on a server and the server can be a component. One or more components may reside within a process and/or thread of execution, and a component may be localized on one computer and/or distributed between two or more computers. Also, these components can execute from various computer readable media having various data structures stored thereon. The components may communicate via local and/or remote processes such as in accordance with a signal having one or more data packets (e.g., data from one component interacting with another component in a local system, distributed system, and/or across

a network such as the Internet with other systems via the signal).

(22) Computer executable components can be stored, for example, at non-transitory, computer readable media including, but not limited to, an ASIC (application specific integrated circuit), CD (compact disc), DVD (digital video disk), ROM (read only memory), floppy disk, hard disk, storage class memory, solid state drive, EEPROM (electrically erasable programmable read only memory), memory stick or any other storage device type, in accordance with the claimed subject matter.

(23) System **100**: FIG. 1A shows an example of a networked operating environment **100** (also referred to as system **100**) used according to one aspect of the present disclosure. As an example, system **100** may include a plurality of storage servers **108A-108N** (may also be referred to as storage server/storage servers/storage controller/storage controllers **108**) executing a storage operating system **114A-114N** (may also be referred to as storage operating system **114** or storage operating systems **114**), a plurality of computing systems **104A-104N** (may also be referred to as server system/server systems **104** or as host system/host systems **104**) that may access storage space provided by a storage-subsystem **112** managed by the storage servers **108** via a connection system **116** such as a local area network (LAN), wide area network (WAN), the Internet and others. The storage-subsystem **112** includes a plurality of storage devices **110A-110N** (may also be referred to as storage device/storage devices/disk/disks **110**) described below in detail. In one aspect, storage devices **110** are ZNS SSDs and are referred to as ZNS SSD or ZNS SSDs **110**, as described below in detail. It is noteworthy that the term “disk” as used herein is intended to mean any storage device/space and not to limit the adaptive aspects to any particular type of storage device, for example, hard disks.

(24) The server systems **104** may communicate with each other via connection system **116**, for example, for working collectively to provide data-access service to user consoles (not shown). Server systems **104** may be computing devices configured to execute applications **106A-106N** (may be referred to as application or applications **106**) over a variety of operating systems, including the UNIX® and Microsoft Windows® operating systems (without derogation of any third-party rights). Application **106** may include an email exchange application, a database application or any other type of application. In another aspect, application **106** may comprise a virtual machine. Applications **106** may utilize storage devices **110** to store and access data.

(25) Server systems **104** generally utilize file-based access protocols when accessing information (in the form of files and directories) over a network attached storage (NAS)-based network. Alternatively, server systems **104** may use block-based access protocols, for example but not limited to, the Small Computer Systems Interface (SCSI) protocol encapsulated over TCP (iSCSI) and SCSI encapsulated over Fibre Channel (FCP) to access storage via a storage area network (SAN).

(26) Server **104** may also execute a virtual machine environment, according to one aspect. In the virtual machine environment, a physical resource is time-shared among a plurality of independently operating processor executable virtual machines (VMs). Each VM may function as a self-contained platform, running its own operating system (OS) and computer executable, application software. The computer executable instructions running in a VM may be collectively referred to herein as “guest software”. In addition, resources available within the VM may be referred to herein as “guest resources”.

(27) The guest software expects to operate as if it were running on a dedicated computer rather than in a VM. That is, the guest software expects to control various events and have access to hardware resources on a physical computing system (may also be referred to as a host platform) which maybe referred to herein as “host hardware resources”. The host hardware resource may include one or more processors, resources resident on the processors (e.g., control registers, caches and others), memory (instructions residing in memory, e.g., descriptor tables), and other resources (e.g., input/output devices, host attached storage, network attached storage or other like storage) that

reside in a physical machine or are coupled to the host platform.

(28) In one aspect, the storage servers **108** use the storage operating system **114** to store and retrieve data from the storage sub-system **112** by accessing the ZNS SSDs **110** via storage device controllers **102A-102N** (may also be referred to as disk controller/disk controllers **110**). Data is stored and accessed using read and write requests that are also referred to as input/output (I/O) requests.

(29) The storage devices **110** may include writable storage device media such as magnetic disks, video tape, optical, DVD, magnetic tape, non-volatile memory devices for example, self-encrypting drives, flash memory devices, ZNS SSDs and any other similar media adapted to store information. The storage devices **110** may be organized as one or more RAID groups. The various aspects disclosed herein are not limited to any particular storage device type or storage device configuration.

(30) In one aspect, ZNS SSDs **110** comply with the NVMe (Non-Volatile Memory Host Controller Interface) zoned namespace (ZNS) specification defined by the NVM Express™ (NVMe™) standard organization. An SSD “zone” as defined by the NVMe ZNS standard is a sequence of blocks that can only be written in a sequential fashion and are overwritten by performing a “Zone Erase” or “Zone Reset operation” per the NVMe specification. A ZNS SSD storage space is exposed as zones. MUs of a ZNS SSD operate independent of each other to store information and are managed by the storage device controller **102**. The zones are configured using the independent MUs, which enables the MUs to operate as individual drives of a RAID group. This enables the storage sub-system **112** to use a single parity ZNS SSD to store parity data and distribute the parity data within each ZNS SSD of a RAID group, as described below in detail.

(31) In one aspect, to facilitate access to ZNS SSDs **110**, the storage operating system **114** “virtualizes” the storage space provided by ZNS SSDs **110**. The storage server **108** can present or export data stored at ZNS SSDs **110** to server systems **104** as a storage volume or one or more qtree sub-volume units. Each storage volume may be configured to store data files (or data containers or data objects), scripts, word processing documents, executable programs, and any other type of structured or unstructured data. From the perspective of the server systems, each volume can appear to be a single drive. However, each volume can represent the storage space in one storage device, an aggregate of some or all the storage space in multiple storage devices, a RAID group, or any other suitable set of storage space.

(32) The storage server **108** may be used to access information to and from ZNS SSDs **110** based on a request generated by server system **104**, a management console (or system) **118** or any other entity. The request may be based on file-based access protocols, for example, the CIFS or the NFS protocol, over TCP/IP. Alternatively, the request may use block-based access protocols, for example, iSCSI or FCP.

(33) As an example, in a typical mode of operation, server system **104** transmits one or more input/output (I/O) commands, such as an NFS or CIFS request, over connection system **116** to the storage server **108**. The storage operating system **114** generates operations to load (retrieve) the requested data from ZNS **110** if it is not resident “in-core,” i.e., at the memory of the storage server. If the information is not in the memory, the storage operating system retrieves a logical volume block number (VBN) that is mapped to disk identifier and disk block number (Disk, DBN). The DBN is accessed from ZNS SSD by the disk controller **102** and loaded in memory for processing by the storage server **108**. Storage server **108** then issues an NFS or CIFS response containing the requested data over the connection system **116** to the respective server system **104**.

(34) In one aspect, storage server **108** may have a distributed architecture, for example, a cluster-based system that may include a separate network module and storage module. Briefly, the network module is used to communicate with host platform server system **104** and management console **118**, while the storage module is used to communicate with the storage subsystem **112**.

(35) The management console **118** executes a management application **117** that is used for

managing and configuring various elements of system **100**. Management console **118** may include one or more computing systems for managing and configuring the various elements.

(36) Parity Protection: Before describing the details of the present disclosure, a brief overview of parity protection in a RAID configuration will be helpful. A parity value for data stored in storage subsystem **112** can be computed by summing (usually modulo 2) data of a particular word size (usually one bit) across a number of similar ZNS SSD holding different data and then storing the results in a parity ZNS SSD. That is, parity may be computed on vectors 1-bit wide, composed of bits in corresponding positions on each ZNS SSD. When computed on vectors 1-bit wide, the parity can be either the computed sum or its complement; these are referred to as even and odd parity, respectively. Addition and subtraction on 1-bit vectors are both equivalent to exclusive-OR (XOR) logical operations. The data is protected against the loss of any one of the ZNS SSDs, or of any portion of the data on any one of the SSDs. If the ZNS SSD storing the parity is lost, the parity can be regenerated from the data or from parity data stored within each ZNS SSD. If one of the ZNS SSD is lost, the data can be regenerated by adding the contents of the surviving ZNS SSDs together and then subtracting the result from the stored parity data.

(37) Typically, storage devices in a RAID configuration are divided into parity groups, each of which comprises one or more data drive and a parity drive. A parity set is a set of blocks, including several data blocks and one parity block, where the parity block is the XOR of all the data blocks. A parity group is a set of drives from which one or more parity sets are selected. The storage space is divided into stripes, with each stripe containing one block from each drive. The blocks of a stripe are usually at the same locations on each drive in the parity group. Within a stripe, all but one block are blocks containing data (“data blocks”) and one block is a block containing parity (“parity block”) computed by the XOR of all the data. The innovative technology described herein uses a single ZNS SSD as a parity drive and stores parity data within each ZNS SSD, as described below in detail.

(38) ZNS SSD RAID Configuration: FIG. 1B illustrates a Hierarchical RAID implementation providing dual parity protection (e.g., RAID6 and RAID-DP) using a single, ZNS SSD **110D** as a parity drive to store parity data, and ZNS SSDs **110A-110C** as data drives storing data. Unlike conventional systems that use two parity drives within a RAID group for providing RAID 6 and RAID-DP type protection, only one parity drive **110D** is used.

(39) Each ZNS SSD **110A-110D** include a plurality of storage blocks identified by disk block numbers (“DBNs”), shown as DBNO-DBNN (e.g., **126A-126N** for ZNS SSD **110A**). The parity drive ZNS SSD **110D** has similar DBNs shown as **128A-128N** for storing parity data. The parity data is computed by XORing data stored at disk blocks in a horizontal stripe with the same DBN of each ZNS SSD data drive (i.e., **110A-110C**). The computed parity is written to the same DBN on the parity drive **110D**. For example, the parity for data stored at the first disk (DBNO) of each ZNS SSD **110A-110C** is stored at the DBNO **128A** of ZNS SSD **110D**. This is referred to as TIER2 RAID for providing RAID protection if a ZNS SSD fails or if a block of a ZNS SSD fails.

(40) Parity is also computed and stored at each ZNS SSD, which is referred to as TIER1 RAID. An example of TIER1 RAID is shown for ZNS SSD **110B** that includes a plurality of MUs **120A-120E**. A plurality of zones is configured for the MUs **120A-120E**, e.g., zones **122A-122C** are based on MU **120A**, while parity zones **124A-124C** are based on MU **120E** to store parity data. The zones within each ZNS SSD are spread uniformly across the MUs. Parity data for TIER1 RAID is computed across zones and stored at the parity zones **124A-124C** within MU **120E**. By grouping zones from independent MUs into a RAID stripe, TIER1 RAID can provide data availability even if a block from one of the zones encounters an uncorrectable read error or an entire MU is inaccessible, as described below in detail.

(41) FIG. 1C illustrates another representation of the innovative dual parity architecture having a single ZNS SSD **110D** within a RAID group to store parity data and storing parity data at each ZNS SSD of the RAID group. A horizontal TIER2 RAID stripe is shown within the rectangle **130**

and the vertical TIER1 RAID stripe is shown within **132**. The vertical TIER1 RAID parity is also shown as L1P0 (**134A-134C**) in ZNS SSDs **110A-110C** and written to disk blocks that are internal to each ZNS SSD, i.e., these hidden disk blocks are not visible to upper software layers (such as TIER2 RAID layer **136** and File System **134** shown in FIG. **1D**, and described below in detail).

(42) Software Architecture: FIG. **1D** shows an example of the innovative software architecture used for implementing the innovative technology of the present disclosure. The architecture includes the file system manager **134** within the storage operating system **114**, described in detail below with respect to FIG. **3**. The TIER2 RAID layer **136** interfaces with the file system manager **134** for processing I/O requests to read and write data. A zone translation layer (ZTL) **138** with a TIER1 RAID layer **140** operate below the TIER2 RAID layer **136** for managing the zones within the ZNS SSDs **110A-110D**. As an example, the total storage capacity of each ZNS SSD is split across physical zones (PZones), e.g., **142** for ZNS SSD **110A** visible only to the TIER1 RAID layer **140**. The PZones are grouped by MUs and each MU may contain a plurality of PZones. The TIER1 RAID layer **140** groups PZones across multiple MUs into a RAID-Zone (“RZone”, e.g., RZone 0 **144** for ZNS SSD **110A**). After the TIER1 RAID layer **140** creates the RZones, the ZTL **138** and upper layers can view each ZNS SSD as a collection of RZones e.g., RZone 0 **146A** and RZone1 **146B** shown for ZNS SSD **110A**.

(43) In one aspect, ZNS SSDs **110A-110D** have defined rules for writing to zones. For example, a zone should be “open: for writing and the writes are sequential with increasing block numbers of the zone. To enable multiple processors to write in parallel, the NVMe ZNS standard allows the ZNS SSDs to provide a Zone Random Write Area (ZRWA) for each available zone. The ZRWA is a buffer within a memory where writes to an open zone are gathered before being written to the PZones. ZRWA enables higher software layers (e.g., file system manager **134** and the TIER2 RAID layer **136**) to issue sequential write commands without the overhead of guaranteeing that the writes arrive in the sequential order at the ZNS SSD. The data from the ZRWA is moved to the ZNS SSD zones via a “commit operation.” An indication for the commit operation is provided by an upper layer software, e.g., the file system manager **134** and/or the TIER2 RAID layer **136**. The commit operation may be explicit or implicit. An explicit commit operation happens when a commit command is sent to the ZNS SSD. An implicit operation commits data to a ZNS SSD zone, when the ZNS SSD receives a write command, which if executed would exceed the size of the ZRWA buffer (i.e., when the ZRWA buffer will reach a threshold value).

(44) Implicit Commit Operations: FIG. **1E** shows an example of using the commit operation in a PZone (e.g., **142**) of a ZNS SSD. Each PZone (e.g., **142**) has a write pointer (WP) (shown as PWP **148**). The location of PWP **148** shows a next writable block within the PZone **142**. When a commit operation is executed, a certain number of data blocks (e.g., **152A/152B**) from the beginning of the ZRWA (shown as PZRWA **150**) are written at the WP **148** of the PZone and the WP **148** is incremented by the number of blocks written. The number of blocks thus written are termed as Commit Granularity (CG) of the PZone. CG is typically a property of the ZNS SSD, shown as an example, as 4 blocks. The size of the ZRWA **150** is a multiple of CG. An implicit commit operation occurs when a software layer sends a write command (shown as **147**) to the ZNS SSD beyond the ZRWA, shown as **152C**. FIG. **1E** shows that the PWP **148** has moved 4 blocks, after the 4 blocks have been committed i.e., transferred to the PZone **142**.

(45) As mentioned above and shown in FIG. **1F**, Tier1 RAID layer **140** constructs virtual RAID Zones (i.e., RZone) by grouping together PZones across multiple MUs, which effectively creates an RZone (e.g., **146**) with an associated ZRWA (shown as RZRWA) **156** and a RZone Write Pointer (shown as RWP) **154**. The example of FIG. **1F** assumes a MU count of 15, which makes the RZRWA size= $15 \times 8 = 120$ blocks and the RCG= $15 \times 4 = 60$ blocks (e.g., **156A/156B**). When a write operation (**158**) exceeds 120 blocks (shown as **156C**), the data is committed from the virtual RZRWA **156** to the SSD. The RWP **154** then slides 60 blocks, as shown in FIG. **1F**. In one aspect, PWP **148** tracks data from PZRWA **150** and RWP **154** track data movement between RZRWA **156**

to RZone **146**. This enables the TIER1 RAID layer to effectively manage data and parity writes, as described below in detail.

(46) PZone/RZone Initialization: FIG. **1G** shows a process **160** for initializing the PZones and RZones by the TIER1 RAID layer **140**, according to one aspect of the present disclosure. The process begins in block **B162**, before a ZNS SSD **110** is made available within the storage subsystem **112**. In block **B164**, the TIER1 RAID layer **140** queries the ZNS SSDs for information regarding the PZones. Each ZNS SSD controller **102** executes firmware instructions out of a ZNS SSD memory. The controller **102** provides information regarding the PZones, which includes a PZone address, size, starting offset value or any other information that can identify the PZone.

(47) In block **B166**, the TIER1 RAID layer **140** groups PZones across independent MUs (e.g., **120A-120E**, FIG. **1B**) to create RZones, e.g., **144** (FIG. **1D**). Thereafter, in block **B168**, the RZones are presented to upper layers, e.g., the TIER2 RAID layer **136**. The TIER2 RAID layer **136** can then present RZones (e.g., **146A**, **146B**, FIG. **1D**) to other layers, e.g., the file system manager **134**. The RZones and the PZones are then used for writing and retrieving data, as well as for storing parity data, as described below in detail. The process then ends in block **B170**.

(48) Parity Generation in TIER1 RAID: FIGS. **2A** and **2B** illustrate parity generation by the TIER1 RAID layer **140**, according to one aspect of the present disclosure. FIG. **2B** shows an example of a process **200** of FIG. **2A** using the ZNS SSD **110B** with independent MUs **120A-120E** (FIG. **1B**). As mentioned above, the upper layers (e.g., the file system manager **134** and the TIER2 RAID layer **136**) only see RZones (e.g., **146A/146B**, FIG. **1D**), hence all write I/Os that are received by the TIER1 RAID layer **140** target an RZone. The TIER1 RAID layer **140** issues child I/Os **224A-224D** to PZones based on a range of blocks that are targeted by the RZone I/O sent by an upper software layer (**134** or **136**). The I/Os **224A-224D** are issued to write data that is temporarily stored at a plurality of I/O buffers **228A-228D** in storage server memory **232**. For example, data associated with I/O **224A** is first written to PZRWA **222A** assigned to the PZone **220A**, before being committed to the PZone **220A**; data for I/O **224B** is written to PZRWA **222B** assigned to the PZone **220B**, before being committed to the PZone **220B**; data for I/O **224C** is written to the PZRWA **222C** assigned to the PZone **220C**, before being committed to the PZone **220C**; and data for I/O **224D** is written to the PZRWA **222D** assigned to the PZone **220D**, before being committed to the PZone **220D**.

(49) The TIER1 RAID layer **140** also computes parity blocks for the parity PZone **220E** corresponding to the targeted RZone. The TIER1 RAID layer **140** issues a parity I/O **226** for computed parity stored at a parity buffer **230**. The parity buffer **230** may be designated within the storage server memory **232** to store parity data. Parity data for I/O **226E** is written to PZRWA **222E** assigned to the PZone **220E**, before being written to the PZone **220E**. The parity data is computed by XORing the data in the I/O buffers **228A-228D**. It is noteworthy that the parity buffer **230** is written to the parity PZone **220E** and committed after all the blocks in a corresponding RZone stripe have been committed to the appropriate PZones (e.g., **220A-220D**). The TIER1 RAID layer **140** assumes that if any RZone I/O targets a block beyond the RZRWA (**156**, FIG. **1F**)+RWP (**154**, FIG. **1F**) then all the I/Os in the data PZones **220A-220D** have been committed. Based on that assumption, the TIER1 RAID layer **140** can write and explicitly commit the parity in the parity buffer **230** to the parity PZone **226**.

(50) Referring now to FIG. **2A**, process **200** begins after a write I/O request is issued by the TIER2 layer **136** (or file system manager **134**). The write I/O provides one or more RZone identifier. The TIER1 layer **140** fetches the I/O request in block **B202**. In block **B204**, the TIER1 layer **140** evaluates the I/O request, determines the size of the data that needs to be written and ascertains the number of blocks that will be required for the I/O request. Based on that determination, the TIER1 RAID layer **140** determines if the I/O request falls within an implicit commit region of the RZone (**156C**, FIG. **1F**). If yes, then in block **B206**, the TIER1 RAID layer **140** determines if all pending write I/Os for the commit region of the RZRWA **156** have been committed to the appropriate

PZones. If not, the I/O is delayed in block B208, until the commit operation is completed.

(51) If the fetched I/O request does not belong to the commit region or if the previous I/O requests for the commit region have been committed, the process moves to block B210, when the parity in parity buffer 230 is updated by XORing the data in the I/O buffers 228A-228D. The TIER1 RAID layer 140 generates child write I/O requests, e.g., 224A-224D, that are sent to the PZRWA 222A-222D and committed to PZones 220A-220D. If there are more I/O requests for the RZone stripe, as determined in block B214, the process reverts back to block B202, otherwise, the TIER1 RAID layer 140 generates a parity I/O 226 that is sent to the PZRWA 222E and committed in block B218. This completes the write I/O request and parity generation by the TIER1 RAID layer 140.

(52) Parity Overwrite: The parity drive may see overwrites to parity blocks when an application sends a write request to write to a partial stripe, data is written to the partial stripe, parity is updated in a RZone of the parity drive 110D, and later, the application sends a new write request to complete the RAID stripe. In this example, the stripe parity is updated by computing the XOR of the new data blocks with the previous parity. This is enabled by using the RZRWA on the RZone of the parity drive 110D because a block in RZRWA is over-writable and an “in-flight parity buffer” can be updated with new data by XOR-ing out the old data in the block and XOR-ing in the new data by which the block is being over-written. The TIER2 RAID layer 136 guarantees that no parity drive write will happen that would result in writing behind the write-pointer 154 for the RZone by providing an indication to the ZTL 138 so that the write pointer 154 can be advanced, described below in detail.

(53) FIG. 2C shows a process 240 for writing to a RZone, according to one aspect of the present disclosure. Process 240 begins when a write request has been received and a next available block is allocated by the file system manager 134 for writing data for the write request. In block 244, the ZTL 138 determines if the block that needs to be rewritten belongs to a certain range identified by the WP 154 (FIG. 1F) and the RZRWA 156 size. The ZTL 138 tracks the WP 154 and is aware of a last written block. If not, then in block B250, the write I/O is sent to the ZTL 138 and handled per the process of FIG. 2A. If yes, then ZTL 138 determines if all the previous blocks for previous one or more write requests, before WP+ZRWA size/2 have been written. If not, then the write I/O is held in block B248 until the previous write requests are complete. If yes, then the write I/O is sent to the ZTL 138 and handled per the process of FIG. 2A.

(54) An example, of process 240 is shown in FIG. 2D that illustrates the I/Os buffered in the storage server memory 232 to ensure that parity drive RZone blocks remain overwritable until a specific TIER2 RAID stripe has been written. FIG. 2D shows the ZNS SSDs 110A-110C with the parity drive 110D. No writes to RAID stripes within commit groups (CGs) 254G, 254H, and 254I with parity 256C are written to the ZNS SSDs until all the writes defined by CGs 254A-254F with parity at 256A/256B have been written. This ensures that all parity updates can be handled sequentially and reduces error risks for parity updates.

(55) Read Operations: To read from a RZone, the ZTL 138 receives a read request and translates logical blocks address (LBAs) for the RZone that are provided by the file system manager 134 to the underlying LBAs of the PZones. The translated LBAs are then used to issue multiple parallel read I/O requests to the ZNS SSDs to obtain data for the read request. An example of the LBA translation is provided below.

(56) For a given raided_zone lba, “rzlba,” a corresponding physical zone LBA, “pzlba” can be determined as follows:

(57) TABLE-US-00001 pzlba (Output) xlate_rzlba_to_pzlba (rzlba (input)) {
rzlba_starting = (rzlba / raided_zone_capacity) *raided_zone_capacity;
rzlba_offset = rzlba – rzlba_starting; stripe_offset = rzlba_offset –
(rzlba_offset / (st_depth * st_width)) * (st_width * st_depth) ;
pzone_index = stripe_offset / st_depth; pzone_start_lba = (rzlba_starting /
(physical_zone_cap * st_width_data)) * (physical_zone_size *

```

st_width_data_parity) ;
st_width_data) ) * st_depth;
pzone_size;
pzone_lba = pzone_lba + (stripe_offset % st_depth) ;
pzone_lba = pzone_lba + pzone_start_lba;
return
pzone_lba;
}

```

(58) The following defines the various parameters of the pseudo code above: **raided_zone**: A set of physical zones grouped together for RAIDed data layout. **physical_zone**: A ZNS zone exposed by a ZNS SSD (e.g. **110A**.) **raided_zone_capacity**: Capacity of a RZone. **physical_zone_capacity**: Capacity of a PZone. **physical_zone_size**: Size of the PZone. **data_zone**: A zone on which data is written. **parity_zone**: A zone holding parity for the data written in the data zones. **st_width_data**: Number of data zones in a stripe. **st_width_data_parity**: Number of zones in a stripe, data and parity. **st_depth**: Number of LBAs in a data zone written before moving to the next data zone. **rzlba**: Externally visible RZone LBA. **pzlba**: PZone LBA.

(59) FIG. 2E shows a process **260** for processing a read request, according to one aspect of the present disclosure. The process begins in block **B262**, after a read I/O request is received by the ZTL **138** from the TIER2 RAID layer **136**. The read request includes a RZone LBA (**rzlba**) and length. In block **B264**, the ZTL **138** translates the **rzlba** into a set of **pzlba** and length pairs. The translation may be executed using the pseudo code described above. The **pzlba** and the length pairs are provided to the TIER1 RAID layer **140** as read I/Os for each PZone LBA in block **B266**. In block **B268**, the TIER1 RAID layer **140** issues read I/O requests to the ZNS SSD that stores the requested data. Once all the requested data has been retrieved, a reply to the read request with the requested data is sent in block **B270**.

(60) Reconstructing Data: FIG. 2F shows a process **276** for reconstructing data when an error is encountered during the read process of FIG. 2E. During a read operation, if there is an error associated with a block and a ZNS SSD indicates a media error, the TIER1 RAID layer **140** reconstructs the data by reading blocks of a stripe associated with the read operation and XORs the blocks with the parity stored at the parity zone, as described above. This prevents propagation of a media error seen from the ZNS SSD to upper layer software because the TIER1 RAID layer **140** can reconstruct the data. The same reconstruction mechanism is used when an independent MU of a ZNS SSD fails and the zones associated with the MU become unavailable. In this instance, the TIER1 RAID layer **140** reconstructs the data for the blocks mapped to those zones during the read operation.

(61) Process **276** begins when a read operation is in progress, as shown in block **B278**. In block **B280**, TIER1 RAID layer **140** determines if all blocks associated with a read request are successfully read, if yes, then the data is returned in block **B282**. If the blocks are not successfully read, then in block **B284**, the TIER1 RAID layer **140** reads each block associated with the read request to identify the block that failed. In block **B286**, for each failed block, other blocks, including the parity block, in the stripe associated with the read request are read. If all the blocks are read, as determined in block **B288**, the failed block is reconstructed by XORing the successfully read data and the parity blocks in block **B290**. The reconstructed data is then returned in block **B292**. If the blocks are not read in block **B288**, then the read operation fails in block **B294** and a failure indication is sent to the file system manager **134**. Storage Operating System: FIG. 3 illustrates a generic example of operating system **114** executed by storage server **108**, according to one aspect of the present disclosure. Storage operating system **114** interfaces with the storage subsystem **112** as described above in detail.

(62) As an example, operating system **114** may include several modules, or “layers”. These layers include a file system manager **134** that keeps track of a directory structure (hierarchy) of the data stored in storage devices and manages read/write operations, i.e., executes read/write operations on disks in response to server system **104** requests.

(63) Operating system **114** may also include a protocol layer **303** and an associated network access

layer **305**, to allow storage server **108** to communicate over a network with other systems, such as server system **104**, and management console **118**. Protocol layer **303** may implement one or more of various higher-level network protocols, such as NFS, CIFS, Hypertext Transfer Protocol (HTTP), TCP/IP and others.

(64) Network access layer **305** may include one or more drivers, which implement one or more lower-level protocols to communicate over the network, such as Ethernet. Interactions between server systems **104** and the storage sub-system **112** are illustrated schematically as a path, which illustrates the flow of data through operating system **114**. The operating system **114** may also include a storage access layer **307** and an associated storage driver layer **309** to communicate with a storage device. The storage access layer **307** may implement a higher-level disk storage protocol, such as TIER2 RAID layer **136**, ZTL **138** and TIER1 RAID layer **140**, while the storage driver layer **309** may implement a lower-level storage device access protocol, such as the NVMe protocol.

(65) It should be noted that the software “path” through the operating system layers described above needed to perform data storage access for a client request may alternatively be implemented in hardware. That is, in an alternate aspect of the disclosure, the storage access request data path may be implemented as logic circuitry embodied within a field programmable gate array (FPGA) or an ASIC. This type of hardware implementation increases the performance of the file service provided by storage server **108**.

(66) As used herein, the term “storage operating system” generally refers to the computer-executable code operable on a computer to perform a storage function that manages data access and may implement data access semantics of a general-purpose operating system. The storage operating system can also be implemented as a microkernel, an application program operating over a general-purpose operating system, such as UNIX® or Windows XP®, or as a general-purpose operating system with configurable functionality, which is configured for storage applications as described herein.

(67) In addition, it will be understood to those skilled in the art that the invention described herein may apply to any type of special-purpose (e.g., file server, filer or storage serving appliance) or general-purpose computer, including a standalone computer or portion thereof, embodied as or including a storage system. Moreover, the teachings of this disclosure can be adapted to a variety of storage system architectures including, but not limited to, a network-attached storage environment, a storage area network and a disk assembly directly-attached to a client or host computer. The term “storage system” should therefore be taken broadly to include such arrangements in addition to any subsystems configured to perform a storage function and associated with other equipment or systems.

(68) Processing System: FIG. 4 is a high-level block diagram showing an example of the architecture of a processing system, at a high level, in which executable instructions as described above can be implemented. The processing system **400** can represent the storage server **108**, the management console **118**, server systems **104**, and others. Note that certain standard and well-known components which are not germane to the present invention are not shown in FIG. 4.

(69) The processing system **400** includes one or more processors **402** and memory **404**, coupled to a bus system **405**. The bus system **405** shown in FIG. 4 is an abstraction that represents any one or more separate physical buses and/or point-to-point connections, connected by appropriate bridges, adapters and/or controllers. The bus system **405**, therefore, may include, for example, a system bus, a Peripheral Component Interconnect (PCI) bus, a HyperTransport or industry standard architecture (ISA) bus, a small computer system interface (SCSI) bus, a universal serial bus (USB), or an Institute of Electrical and Electronics Engineers (IEEE) standard **1394** bus (sometimes referred to as “Firewire”).

(70) The processors **402** are the central processing units (CPUs) of the processing system **400** and, thus, control its overall operation. In certain aspects, the processors **402** accomplish this by executing programmable instructions stored in memory **404**. A processor **402** may be, or may

include, one or more programmable general-purpose or special-purpose microprocessors, digital signal processors (DSPs), programmable controllers, application specific integrated circuits (ASICs), programmable logic devices (PLDs), or the like, or a combination of such devices. (71) Memory **404** represents any form of random-access memory (RAM), read-only memory (ROM), flash memory, or the like, or a combination of such devices. Memory **404** includes the main memory of the processing system **400**. Instructions **406** which implements techniques introduced above may reside in and may be executed (by processors **402**) from memory **404**. For example, instructions **406** may include code for executing the process blocks of FIGS. **1G**, **2A**, **2C** and **2E-2F**.

(72) Also connected to the processors **402** through the bus system **405** are one or more internal mass storage devices **410**, and a network adapter **412**. Internal mass storage devices **410** may be or may include any conventional medium for storing large volumes of data in a non-volatile manner, such as one or more magnetic or optical based disks. The network adapter **412** provides the processing system **400** with the ability to communicate with remote devices (e.g., storage servers) over a network and may be, for example, an Ethernet adapter, a FC adapter, or the like. The processing system **400** also includes one or more input/output (I/O) devices **408** coupled to the bus system **405**. The I/O devices **408** may include, for example, a display device, a keyboard, a mouse, etc.

(73) Cloud Computing: The system and techniques described above are applicable and especially useful in the cloud computing environment where storage at ZNS **110** is presented and shared across different platforms. Cloud computing means computing capability that provides an abstraction between the computing resource and its underlying technical architecture (e.g., servers, storage, networks), enabling convenient, on-demand network access to a shared pool of configurable computing resources that may be rapidly provisioned and released with minimal management effort or service provider interaction. The term “cloud” is intended to refer to a network, for example, the Internet and cloud computing allows shared resources, for example, software and information to be available, on-demand, like a public utility.

(74) Typical cloud computing providers deliver common business applications online which are accessed from another web service or software like a web browser, while the software and data are stored remotely on servers. The cloud computing architecture uses a layered approach for providing application services. A first layer is an application layer that is executed at client computers. In this example, the application allows a client to access storage via a cloud.

(75) After the application layer is a cloud platform and cloud infrastructure, followed by a “server” layer that includes hardware and computer software designed for cloud specific services. The storage systems described above may be a part of the server layer for providing storage services. Details regarding these layers are not germane to the inventive aspects.

(76) Thus, a method and apparatus for protecting data using ZNS SSDs within system **100** have been described. Note that references throughout this specification to “one aspect” or “an aspect” mean that a particular feature, structure or characteristic described in connection with the aspect is included in at least one aspect of the present invention. Therefore, it is emphasized and should be appreciated that two or more references to “an aspect” or “one aspect” or “an alternative aspect” in various portions of this specification are not necessarily all referring to the same aspect. Furthermore, the particular features, structures or characteristics being referred to may be combined as suitable in one or more aspects of the present disclosure, as will be recognized by those of ordinary skill in the art.

(77) While the present disclosure is described above with respect to what is currently considered its preferred aspects, it is to be understood that the disclosure is not limited to that described above. To the contrary, the disclosure is intended to cover various modifications and equivalent arrangements within the spirit and scope of the appended claims.

Claims

1. A method executed by one or more processor, comprising; presenting a plurality of physical zones (PZones) to a first tier RAID (redundant array of independent disks) layer, and a plurality of RAID zones (RZones) to a second tier RAID layer, the PZones based on storage space of a plurality of zoned solid-state drives (ZNS SSDs) of a storage subsystem and each RZone having a plurality of PZones; translating, by a zone translation layer, a logical block address (“LBA”) of a RZone specified by a read request into a LBA of a corresponding PZone that stores data for the read request; and utilizing, by the first tier RAID layer, the LBA of the PZone to retrieve data for the read request.
2. The method of claim 1, further comprising: in response to an error in reading a portion of the data associated with the read request, reconstructing the portion of the data from parity information corresponding to the portion of the data stored in a parity ZNS SSD.
3. The method of claim 1, further comprising: in response to an error in reading a portion of the data associated with the read request, reconstructing the portion of the data from parity information corresponding to the portion of the data stored in a parity zone of one of the ZNS SSDs, when the parity information is unavailable from a parity ZNS SSD.
4. The method of claim 1, further comprising: generating the plurality of PZones across a plurality of independent media units of each ZNS SSD.
5. The method of claim 1, further comprising: interfacing, by the ZTL, with the first tier RAID layer and the second tier RAID layer to process the read request.
6. The method of claim 1, further comprising: receiving, by the second tier RAID layer, the read request and providing, by the second tier RAID layer, the LBA of the RZone to the ZTL.
7. The method of claim 1, further comprising: providing, by the ZTL, the LBA of the PZone to the first tier RAID layer for retrieving the data.
8. A non-transitory, machine-readable storage medium having stored thereon instructions for performing a method, comprising machine executable code which when executed by at least one machine, causes the machine to: present a plurality of physical zones (PZones) to a first tier RAID (redundant array of independent disks) layer, and a plurality of RAID zones (RZones) to a second tier RAID layer, the PZones based on storage space of a plurality of zoned solid-state drives (ZNS SSDs) of a storage subsystem and each RZone having a plurality of PZones; translate, by a zone translation layer, a logical block address (“LBA”) of a RZone specified by a read request into a LBA of a corresponding PZone that stores data for the read request; and utilize, by the first tier RAID layer, the LBA of the PZone to retrieve data for the read request.
9. The non-transitory, machine-readable storage medium of claim 8, wherein the machine executable code which when executed by at least one machine, further causes the machine to: in response to an error in reading a portion of the data associated with the read request, reconstruct the portion of the data from parity information corresponding to the portion of the data stored in a parity ZNS SSD.
10. The non-transitory, machine-readable storage medium of claim 8, wherein the machine executable code which when executed by at least one machine, further causes the machine to: in response to an error in reading a portion of the data associated with the read request, reconstructing the portion of the data from parity information corresponding to the portion of the data stored in a parity zone of one of the ZNS SSDs, when the parity information is unavailable from a parity ZNS SSD.
11. The non-transitory, machine-readable storage medium of claim 8, wherein the machine executable code which when executed by at least one machine, further causes the machine to: generate the plurality of PZones across a plurality of independent media units of each ZNS SSD.
12. The non-transitory, machine-readable storage medium of claim 8, wherein the machine

- executable code which when executed by at least one machine, further causes the machine to: interface, by the ZTL, with the first tier RAID layer and the second tier RAID layer to process the read request.
13. The non-transitory, machine-readable storage medium of claim 8, wherein the machine executable code which when executed by at least one machine, further causes the machine to: receive, by the second tier RAID layer, the read request and providing, by the second tier RAID layer, the LBA of the RZone to the ZTL.
14. A system, comprising: a memory containing machine readable medium comprising machine executable code having stored thereon instructions; and a processor module coupled to the memory, the processor module configured to execute the machine executable code to: present a plurality of physical zones (PZones) to a first tier RAID (redundant array of independent disks) layer, and a plurality of RAID zones (RZones) to a second tier RAID layer, the PZones based on storage space of a plurality of zoned solid-state drives (ZNS SSDs) of a storage subsystem and each RZone having a plurality of PZones; translate, by a zone translation layer, a logical block address (“LBA”) of a RZone specified by a read request into a LBA of a corresponding PZone that stores data for the read request; and utilize, by the first tier RAID layer, the LBA of the PZone to retrieve data for the read request.
15. The system of claim 14, wherein the machine executable code further causes to: in response to an error in reading a portion of the data associated with the read request, reconstruct the portion of the data from parity information corresponding to the portion of the data stored in a parity ZNS SSD.
16. The system of claim 14, wherein the machine executable code further causes to: in response to an error in reading a portion of the data associated with the read request, reconstructing the portion of the data from parity information corresponding to the portion of the data stored in a parity zone of one of the ZNS SSDs, when the parity information is unavailable from a parity ZNS SSD.
17. The system of claim 14, wherein the machine executable code further causes to: generate the plurality of PZones across a plurality of independent media units of each ZNS SSD.
18. The system of claim 17, wherein the machine executable code further causes to: interface, by the ZTL, with the first tier RAID layer and the second tier RAID layer to process the read request.
19. The system of claim 14, wherein the machine executable code further causes to: receive, by the second tier RAID layer, the read request and providing, by the second tier RAID layer, the LBA of the RZone to the ZTL.
20. The system of claim 19, wherein the machine executable code further causes to: provide, by the ZTL, the LBA of the PZone to the first tier RAID layer for retrieving the data.
-