

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent
Kind Code
Date of Patent
Inventor(s)

12405859
B2
September 02, 2025
Agarwal; Shobhit et al.

Standbys for continuous data protection-enabled objects

Abstract

A stream of data changes to content of an object associated with a source system is received. A hot standby version of at least a portion of the object maintained at a recovery system is updated by streaming the received stream of data changes to the recovery system. It is determined that a gap exists in the stream of data changes. In response to determining that the gap exists in the stream of data changes, a reference snapshot is requested from the source system. A current state of the hot standby version of at least the portion of the object maintained at the recovery system is caused to be updated to a state of the portion of the object associated with the reference snapshot.

Inventors: Agarwal; Shobhit (San Jose, CA), Manjunath; Chinmaya (San Jose, CA), Nerella; Kishan Venkata Sravan (San Jose, CA), Mehra; Himanshu (Mountain View, CA)

Applicant: Cohesity, Inc. (San Jose, CA)

Family ID: 81848072

Assignee: Cohesity, Inc. (Santa Clara, CA)

Appl. No.: 18/408314

Filed: January 09, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240143450 A1	May. 02, 2024

Related U.S. Application Data

continuation parent-doc US 17540630 20211202 US 11914480 child-doc US 18408314
us-provisional-application US 63122722 20201208

Publication Classification

Int. Cl.: G06F16/00 (20190101); G06F11/14 (20060101); G06F16/11 (20190101); G06F16/17 (20190101); G06F16/901 (20190101)

U.S. Cl.:

CPC **G06F11/1458** (20130101); **G06F16/128** (20190101); **G06F16/1734** (20190101); **G06F16/9027** (20190101); G06F11/1461 (20130101); G06F2201/84 (20130101)

Field of Classification Search

CPC: G06F (11/1458); G06F (2201/84)

USPC: 707/654

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7421648	12/2007	Davis	N/A	N/A
7437764	12/2007	Sobel	N/A	N/A
8020037	12/2010	Schwartz	N/A	N/A
8086585	12/2010	Brashers	N/A	N/A
8112661	12/2011	La France	N/A	N/A
8190583	12/2011	Shekar	N/A	N/A
8312471	12/2011	Davis	N/A	N/A
8364648	12/2012	Sim-Tang	N/A	N/A
8364652	12/2012	Vijayan	707/694	G06F 16/9014
8364920	12/2012	Parkison	711/161	G06F 11/1453
8607342	12/2012	Liao	N/A	N/A
9268689	12/2015	Chen	N/A	N/A
9304864	12/2015	Bushman	N/A	N/A
9311190	12/2015	Bushman	N/A	N/A
9361185	12/2015	Bushman	N/A	N/A
9396202	12/2015	Drobychev	N/A	G06F 16/178
9471441	12/2015	Lyadvinsky	N/A	N/A
9489230	12/2015	Patwardhan	N/A	N/A
9594514	12/2016	Bono	N/A	N/A
9621428	12/2016	Lev	N/A	N/A
9983812	12/2017	Don	N/A	N/A
10037223	12/2017	Park	N/A	N/A
10089148	12/2017	Blitzer	N/A	G06F 9/5027
10162528	12/2017	Sancheti	N/A	N/A
10169077	12/2018	Sigl, Sr.	N/A	N/A
10175896	12/2018	Battaje	N/A	N/A
10275321	12/2018	Bajaj	N/A	N/A
10496497	12/2018	Yadav	N/A	N/A
10503612	12/2018	Wang	N/A	N/A
10545776	12/2019	Kowalski	N/A	N/A
10628270	12/2019	Bajaj	N/A	G06F 11/1471
10877928	12/2019	Nagrале	N/A	N/A
10896097	12/2020	Purcell	N/A	N/A
11036594	12/2020	Shats	N/A	N/A
11042452	12/2020	McNutt	N/A	G06F 11/1466
11086545	12/2020	Dayal	N/A	N/A
11176154	12/2020	Dasgupta	N/A	N/A
11275834	12/2021	Offer	N/A	N/A
2003/0033344	12/2002	Abbott	N/A	N/A

2004/0250033	12/2003	Prahlad	N/A	N/A
2005/0055520	12/2004	Kappler	N/A	N/A
2005/0243817	12/2004	Wrenn	370/389	H04L 69/329
2006/0069861	12/2005	Amano	N/A	N/A
2006/0182255	12/2005	Luck, Jr.	N/A	N/A
2007/0153675	12/2006	Baglin	N/A	N/A
2008/0208926	12/2007	Smoot	N/A	N/A
2009/0089657	12/2008	Davis	N/A	N/A
2009/0171707	12/2008	Bobak	N/A	N/A
2009/0313503	12/2008	Atluri	N/A	N/A
2009/0327445	12/2008	Van Rietschote	N/A	N/A
2010/0031170	12/2009	Carullo	N/A	N/A
2010/0070725	12/2009	Prahlad	N/A	N/A
2010/0106933	12/2009	Kamila	N/A	N/A
2010/0122248	12/2009	Robinson	N/A	N/A
2011/0022879	12/2010	Chavda	N/A	N/A
2011/0106776	12/2010	Vik	N/A	N/A
2011/0107246	12/2010	Vik	N/A	N/A
2012/0203742	12/2011	Goodman	N/A	N/A
2013/0006943	12/2012	Chavda	N/A	N/A
2013/0179481	12/2012	Halevy	N/A	N/A
2013/0191347	12/2012	Bensinger	N/A	N/A
2013/0219135	12/2012	Knowles	N/A	N/A
2013/0227558	12/2012	Du	N/A	N/A
2013/0232480	12/2012	Winterfeldt	N/A	N/A
2013/0232497	12/2012	Jalagam	N/A	N/A
2013/0254402	12/2012	Anand	N/A	N/A
2013/0262387	12/2012	Varadharajan	N/A	N/A
2013/0322335	12/2012	Smith	N/A	N/A
2014/0040206	12/2013	Ramakrishnan	N/A	N/A
2014/0052692	12/2013	Zhang	N/A	N/A
2014/0059306	12/2013	Bender	N/A	N/A
2014/0165060	12/2013	Muller	N/A	N/A
2014/0297588	12/2013	Babashetty	N/A	N/A
2014/0359229	12/2013	Cota-Robles	N/A	N/A
2014/0372553	12/2013	Blackburn	N/A	N/A
2015/0193487	12/2014	Demidov	N/A	N/A
2015/0254150	12/2014	Gordon	N/A	N/A
2015/0269032	12/2014	Muthyala	707/649	G06F 16/178
2015/0278046	12/2014	Zellermayer	N/A	N/A
2015/0347242	12/2014	Martos	N/A	N/A
2015/0363270	12/2014	Hammer	N/A	N/A
2015/0370502	12/2014	Aron	N/A	N/A
2015/0378765	12/2014	Singh	N/A	N/A
2016/0004450	12/2015	Lakshman	N/A	N/A
2016/0034356	12/2015	Aron	N/A	N/A
2016/0048408	12/2015	Madhu	N/A	N/A
2016/0070714	12/2015	D Sa	N/A	N/A
2016/0085575	12/2015	Dornemann	N/A	N/A
2016/0085636	12/2015	Dornemann	N/A	N/A
2016/0098323	12/2015	Mutha	N/A	N/A
2016/0125059	12/2015	Jain	N/A	N/A
2016/0162378	12/2015	Garlapati	N/A	N/A
2016/0188898	12/2015	Karinta	N/A	N/A

2016/0203060	12/2015	Singh	N/A	N/A
2016/0232061	12/2015	Gaschler	N/A	N/A
2016/0321339	12/2015	Tekade	N/A	N/A
2016/0357640	12/2015	Bushman	N/A	N/A
2016/0357641	12/2015	Bushman	N/A	N/A
2016/0357769	12/2015	Bushman	N/A	N/A
2017/0031613	12/2016	Lee	N/A	N/A
2017/0031622	12/2016	Nagarajan	N/A	N/A
2017/0060710	12/2016	Ramani	N/A	N/A
2017/0060884	12/2016	Goodman	N/A	N/A
2017/0123935	12/2016	Pandit	N/A	N/A
2017/0168903	12/2016	Dornemann	N/A	N/A
2017/0185488	12/2016	Kumarasamy	N/A	G06F 11/2038
2017/0185491	12/2016	Hajare	N/A	N/A
2017/0185729	12/2016	Boray	N/A	N/A
2017/0193116	12/2016	Wong	N/A	N/A
2017/0206212	12/2016	Srivilliputtur Mannarswamy	N/A	N/A
2017/0212680	12/2016	Waghulde	N/A	N/A
2017/0262347	12/2016	Dornemann	N/A	N/A
2017/0337109	12/2016	Ramu	N/A	N/A
2018/0004437	12/2017	Battaje	N/A	N/A
2018/0004764	12/2017	Sudarsanam	N/A	N/A
2018/0060106	12/2017	Madtha	N/A	N/A
2018/0060187	12/2017	Chavda	N/A	N/A
2018/0081766	12/2017	Ghuge	N/A	N/A
2018/0081902	12/2017	McKenzie	N/A	N/A
2018/0088973	12/2017	Subhraveti	N/A	N/A
2018/0095846	12/2017	Sanakkayala	N/A	N/A
2018/0113625	12/2017	Sancheti	N/A	N/A
2018/0196820	12/2017	Kremer	N/A	N/A
2018/0212896	12/2017	Chang	N/A	N/A
2018/0253414	12/2017	Hailpern	N/A	N/A
2018/0293374	12/2017	Chen	N/A	N/A
2018/0316577	12/2017	Freeman	N/A	N/A
2018/0329637	12/2017	Battaje	N/A	N/A
2019/0065277	12/2018	Raikov	N/A	N/A
2019/0073276	12/2018	Yuen	N/A	N/A
2019/0108266	12/2018	Manvar	N/A	N/A
2019/0129799	12/2018	Kumarasamy	N/A	N/A
2019/0132203	12/2018	Wince	N/A	N/A
2019/0197020	12/2018	Yap	N/A	N/A
2019/0215358	12/2018	Kobayashi	N/A	N/A
2019/0220198	12/2018	Kashi Visvanathan	N/A	N/A
2019/0228097	12/2018	Kassa	N/A	N/A
2019/0278662	12/2018	Nagrале	N/A	N/A
2019/0278663	12/2018	Mehta	N/A	N/A
2020/0026538	12/2019	Cui	N/A	N/A
2020/0034254	12/2019	Natanzon	N/A	N/A
2020/0057567	12/2019	Hutcheson	N/A	N/A
2020/0057669	12/2019	Hutcheson	N/A	N/A
2020/0110755	12/2019	Waldman	N/A	N/A
2020/0142865	12/2019	Manjunath	N/A	N/A
2020/0159625	12/2019	Hutcheson	N/A	N/A
2020/0167238	12/2019	Killamsetti	N/A	N/A

2020/0183794	12/2019	Dwarampudi	N/A	N/A
2020/0233571	12/2019	Yuravlivker	N/A	N/A
2020/0278274	12/2019	Shetty	N/A	N/A
2020/0285449	12/2019	McIntosh	N/A	N/A
2020/0394072	12/2019	Sreekantaswamy	N/A	N/A
2021/0056203	12/2020	Qiao	N/A	N/A
2021/0081087	12/2020	Wayne	N/A	N/A
2021/0081432	12/2020	Grunwald	N/A	G06F 11/1471
2021/0103556	12/2020	Nagrale	N/A	N/A
2021/0117293	12/2020	Luo	N/A	N/A
2021/0232579	12/2020	Schechter	N/A	N/A
2021/0318851	12/2020	Sahu	N/A	N/A
2021/0382771	12/2020	Schniebel	N/A	N/A
2021/0385254	12/2020	Pettit	N/A	N/A
2022/0035714	12/2021	Schultz	N/A	N/A
2022/0138057	12/2021	Juniwal	N/A	N/A
2022/0179748	12/2021	Agarwal et al.	N/A	N/A

OTHER PUBLICATIONS

Prosecution History from U.S. Appl. No. 17/540,630, now issued U.S. Pat. No. 11,914,480, dated Jun. 23, 2023 through Nov. 6, 2023, 55 pp. cited by applicant

“Backup Solution Guide”—Synology [https://download.synology.com/download/www-res/brochure/backup_solution_guide_en-global .pdf](https://download.synology.com/download/www-res/brochure/backup_solution_guide_en-global.pdf) (Year: 2019). cited by applicant

“Recovering File from an Amazon EBS Volume Backup”—Josh Rad, AWS, Feb. 1, 2019

<https://aws.amazon.com/blogs/compute/recovering-files-from-an-amazon-efs-volume-backup/> (Year: 2019). cited by applicant

Actifio. “Getting Started with Actifio VDP.” Sep. 23, 2020.

<https://web.archive.org/web/20200923181125/https://docs.actifio.com/10.0/PDFs/Introducing.pdf> (Year: 2020). cited by applicant

C. Grace. “Site Recovery Manager Technical Overview.” Dec. 1, 2020.

<https://web.archive.org/web/20201201181602/https://core.vmware.com/resource/site-recovery-manager-technical-overview> (Year: 2020). cited by applicant

Cloud Endure. “Cloud Endure Documentation.” Dec. 1, 2020.

<https://web.archive.org/web/20201201022045/https://docs.cloudendure.com/CloudEndure%20Documentation.htm> (Year: 2020). cited by applicant

Cohesity, Cohesity Data Protection White Paper, 2016, Cohesity, pp. 1-12 (Year: 2016). cited by applicant

Gaetan Castlelein, Cohesity SnapFS and SnapTree, Aug. 9, 2017, Cohesity, pp. 1-4 (Year: 2017). cited by applicant

M. Chuang. “Announcing VMware Cloud Disaster Recovery.” Sep. 29, 2020.

<https://web.archive.org/web/20201102133037/https://blogs.vmware.com/virtualblocks/2020/09/29/announcing-vmware-cloud-disaster-recovery/> (Year: 2020). cited by applicant

M. McLaughlin. “VMware Cloud Disaster Recovery is Now Available.” Oct. 20, 2020.

<https://web.archive.org/web/20201103021801/https://blogs.vmware.com/virtualblocks/2020/10/20/vmware-cloud-disaster-recovery-is-now-available/> (Year: 2020). cited by applicant

Red Hat. “Red Hat Virtualization 4.3 Disaster Recovery Guide.” Jul. 17, 2019.

https://web.archive.org/web/20190717013417/https://access.redhat.com/documentation/en-us/red_hat_virtualization/4.3/html/disaster_recovery_guide/index (Year: 2019). cited by applicant

Red Hat. “Red Hat Virtualization 4.3 Product Guide.” Jul. 17, 2019.

https://web.archive.org/web/20190717013254/https://access.redhat.com/documentation/en-us/red_hat_virtualization/4.3/html/product_guide/index (Year: 2019). cited by applicant

VMware. “Site Recovery Manager Administration.” May 31, 2019. <https://docs.vmware.com/en/Site-Recovery-Manager/8.5/srm-admin-8-5.pdf> (Year: 2019). cited by applicant

VMware. “Site Recovery Manager Evaluation Guide.” Oct. 19, 2020.

<https://web.archive.org/web/20201019155135/https://core.vmware.com/resource/site-recovery-manager->

Primary Examiner: Pyo; Monica M

Attorney, Agent or Firm: Shumaker & Sieffert, P.A.

Background/Summary

CROSS REFERENCE TO OTHER APPLICATIONS (1) This application is a continuation of U.S. patent application Ser. No. 17/540,630 entitled STANDBYS FOR CONTINUOUS DATA PROTECTION-ENABLED OBJECTS filed Dec. 2, 2021, which claims priority to U.S. Provisional Patent Application No. 63/122,722 entitled STANDBYS FOR CONTINUOUS DATA PROTECTION-ENABLED OBJECTS filed Dec. 8, 2020, each of which is incorporated herein by reference for all purposes.

BACKGROUND OF THE INVENTION

(1) A data management provider may provide disaster recovery services for one or more objects associated with an entity. Examples of objects include files, databases, virtual machine (VM), applications, containers, pods, etc. An entity may correspond to a user, an individual, an enterprise, a government, a company, an organization, etc. A source system may host the one or more objects associated with the entity.

(2) A user associated with the entity may specify a data protection intent that indicates a primary site, a disaster recovery site, a specified recovery point objective (RPO) and a specified recovery time objective (RTO) for one or more of the objects associated with the entity. RPO is a measure of the acceptable amount of data loss for the one or more objects after a disruption of the source system. RTO is a measure of how fast the one or more objects need to be back online after the disruption. The combination of the specified RPO and the specified RTO may require a corresponding hot standby for each of the one or more objects to be generated and maintained. A hot standby is a replica of an object associated with the source system that is up-to date and ready to use within the specified RTO.

(3) Other systems may generate a corresponding hot standby for each of the one or objects at a disaster recovery site by sending real-time updates from a primary site source system to a disaster recovery site system. However, sending updates directly from the primary site source system to the disaster recovery site system may reduce the performance of the primary site source system because sending the real-time updates may use resources (e.g., processing resources, network resources) that the source system needs to execute the one or more objects.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) Various embodiments of the invention are disclosed in the following detailed description and the accompanying drawings.

(2) FIG. 1A is a block diagram illustrating a system for generating hot standbys for continuous data protection-enabled objects in accordance with some embodiments.

(3) FIG. 1B is a block diagram of a system for providing continuous data protection in accordance with some embodiments.

(4) FIG. 1C is a block diagram illustrating an example of a source system in accordance with some embodiments.

(5) FIG. 1D is a block diagram illustrating an example of a backup system in accordance with some embodiments.

(6) FIG. 2 is an example of a continuous data protection timeline in accordance with some embodiments.

(7) FIG. 3A is a block diagram illustrating an embodiment of a tree data structure.

(8) FIG. 3B is a block diagram illustrating an embodiment of a cloned snapshot tree.

(9) FIG. 3C is a block diagram illustrating an embodiment of modifying a cloned snapshot tree.

- (10) FIG. 3D is a block diagram illustrating an embodiment of a modified cloned snapshot tree.
- (11) FIG. 4A is a block diagram illustrating an embodiment of a tree data structure.
- (12) FIG. 4B is a block diagram illustrating an embodiment of a cloned metadata structure.
- (13) FIG. 4C is a block diagram illustrating an embodiment of modifying a cloned metadata structure.
- (14) FIG. 4D is a block diagram illustrating an embodiment of a modified cloned metadata structure.
- (15) FIG. 5 is a flow chart of a process for providing continuous data protection in accordance with some embodiments.
- (16) FIG. 6 is a flow chart illustrating a process of providing a reference restoration point in accordance with some embodiments.
- (17) FIG. 7 is a flow chart illustrating a process for providing a reference restoration point in accordance with some embodiments.
- (18) FIG. 8 is a flow chart illustrating a process for providing continuous data protection in accordance with some embodiments.
- (19) FIG. 9 is a flow chart illustrating a process for implementing a disaster recovery plan in accordance with some embodiments.
- (20) FIG. 10 is a flow chart illustrating a process for generating a hot standby for a continuous data protection-enabled object in accordance with some embodiments.
- (21) FIG. 11 is a flow chart illustrating a process for generating and maintaining a hot standby in accordance with some embodiments.
- (22) FIG. 12 is a flow chart illustrating a process for updating a hot standby in accordance with some embodiments.
- (23) FIG. 13 is a flow chart illustrating a process for updating a hot standby in accordance with some embodiments.
- (24) FIG. 14 is a flow chart illustrating a process for updating a hot standby in accordance with some embodiments.

DETAILED DESCRIPTION

- (25) A technique to generate hot standbys for continuous data protection-enabled objects is disclosed herein. A source system may host one or more objects. In the event of a failure of an object hosted on the source system or a failure of the source system itself, the technique disclosed herein minimizes the amount of downtime associated with the object. Other systems may restore the object from a backup, which may take a long time to perform. Other systems may generate a hot standby for an object at a disaster recovery site by sending real-time updates from a primary site source system to a disaster recovery site system. However, this may reduce the performance of the primary site source system because sending the real-time updates may use resources (e.g., processing resources, network resources) that the source system needs to execute the one or more objects.
- (26) A user associated with an entity may specify, via a graphical user interface (GUI) associated with a user device, a data protection intent for one or more objects hosted by the source system. The data protection intent includes an indication of a data pool and one or more data management services (e.g., backup, restore, disaster recovery, replication, migration, analytics, compliance, security, etc.) that are to be applied to the data pool.
- (27) A data pool is a description of one or more objects to be included in the data pool. When the one or more data management services are performed, the one or more data management services are performed with respect to the one or objects specified by the data pool. The description may indirectly and/or directly identify the one or more objects. An object may be associated with one or more tags, labels, features, and/or prefixes (e.g., a naming convention). The one or more objects included in the data pool may change over time. For example, the description may be assessed immediately (e.g., as soon as the intent-based data management plan is validated, according a schedule, each time a service is performed, etc.). A user may specify the one or more objects to be included in the data pool when the one or more data management services are performed based on a tag, label, feature, and/or prefix associated with an object.
- (28) The one or more objects included in the data pool may change over time. Objects may be added to or removed from a data pool. For example, a first set of objects may be included in the data pool according to the data pool description at a first point in time when a first instance of a first management service is performed. The first management service may be performed for one or more objects having a first tag. After the first management service is performed, the first tag may be removed from one or more objects included

in the first set of objects. A second set of objects may be included in the data pool according to the data pool description at a second point in time when a second instance of the first management service is performed. The second set of objects may include the one or more objects that were included in the first set of objects that still have the first tag, but exclude the one or more objects that were included in the first set of objects, but no longer have the first tag.

(29) The data protection intent specification may include a disaster recovery plan, i.e., a data management service, for the one or more objects to be included in the data pool when the disaster recovery plan is performed. The disaster recovery plan may indicate a primary site associated with each of the one or more objects, a disaster recovery site associated with each of the one or more objects, a specified RPO for the one or more objects, and/or a specified RTO for the one or more objects. The data protection intent specification may be provided from the user device to a data management plan orchestrator.

(30) The data management plan orchestrator may analyze the data protection intent specification to determine configurations for one or more data storage services needed to achieve the data protection intent specification. The one or more data storage services may include protection jobs, sync-replication, and/or hot-standby. The data protection intent specification analysis may identify one or more types of objects specified to be included in the data pool when the disaster recovery plan is performed. For example, the data management plan orchestrator may identify, based on tags, labels, and/or features associated with a particular object (e.g., file size, creation date, ownership, last modification date, etc.), whether the particular object is a high input/output(IO) object (e.g., an object having a number of IOs/min greater than or equal to an IO threshold), whether the particular object is a low IO object (e.g., an object having a number of IOs/min less than or equal to the IO threshold), performance requirements associated with the particular object, etc.

(31) The data management plan orchestrator may determine protection jobs for the one or more objects to be included in the data pool when the disaster recovery plan is performed based on a specified RPO included in the data protection intent specification and the data protection intent specification analysis. In some embodiments, the data management plan orchestrator determines that a scheduled backup (e.g., every hour, every six hours, daily, etc.) is sufficient to achieve the specified RPO for an object. In some embodiments, the data management plan orchestrator determines that continuous data protection (CDP) needs to be enabled to achieve the specified RPO for an object. When enabling CDP, the data management plan orchestrator may increase a rate at which data associated with an object is backed up from a source system hosting the object to a backup system without reducing a performance associated with the object. CDP may also reduce the amount of data associated with an object that is lost as a result of the source system going offline.

(32) The data management plan orchestrator may automatically enable CDP for one or more objects to be included in the data pool when the disaster recovery plan is performed that are determined to need CDP to achieve a specified RPO by providing to a backup system associated with the source system one or more commands to enable CDP for the one or more objects determined to need CDP to achieve the specified RPO. In response to receiving the one or more commands, the backup system may provide to the source system a request to perform a corresponding reference snapshot and a corresponding write filter for each of the one or more objects determined to need CDP to achieve the specified RPO and enable on the backup system a CDP service to handle write IOs. In some embodiments, the source system includes a corresponding write filter for each of the one or more objects and the one or more commands enable the corresponding write filters. In response to receiving the request to perform a corresponding reference snapshot, the source system may perform a full backup snapshot and include all of the object data associated with the one or more objects determined to need CDP to achieve the specified RPO. The backup system may ingest the data included in the reference snapshot and generate a tree data structure for the one or more objects determined to need CDP to achieve the specified RPO. In some embodiments, the tree data structure includes a snapshot tree for an object and one or more metadata structures that corresponding to one or more object files associated with the object. In some embodiments, the tree data structure is a snapshot tree for some or all of the objects hosted on the source system and a corresponding metadata structure for each of the one or more objects hosted by the source system that are to be included in the data pool when the disaster recovery plan is performed. The tree data structure corresponding to the reference snapshot may be a reference restoration point for an object and may be referred to as a “view corresponding to a reference restoration point.”

(33) In response to receiving a write filter, the source system may attach the corresponding write filter to each of the one or more objects determined to need CDP to achieve the specified RPO. A write filter may be configured to intercept all IOs of an associated object, filter the read IOs, and provide a stream of the intercepted write IOs to the backup system. The write filter may be configured to attach a sequence number to each intercepted write IO. In some embodiments, the backup system is configured to provide an IO filter daemon to the source system. The IO filter daemon of the source system may be configured to receive the intercepted write IOs from each of the write filters of the source system and to provide the intercepted write IOs from each of the write filters of the source system to the backup system.

(34) The backup system may include one or more CDP services that are configured to receive one or more intercepted write IOs from the IO filter daemon of the source system. Each cluster node of the backup system may have a corresponding CDP service. The IO filter daemon may select one of the one or more CDP services and provide the intercepted write IOs associated with a particular object to the selected CDP service. A CDP service may maintain a corresponding log for each of the CDP-enabled objects to store the one or more intercepted write IOs received from the IO filter daemon of the source system. Each entry of a log may store the metadata associated with an intercepted write IO. For example, the metadata associated with an intercepted write IO may include at least one of the following: sequence number, timestamp, IO offset, IO length, IO data bytes, VM universally unique identifier (UUID), VM disk (VMDK) full path on a host at which the VM is running, etc.

(35) The data management plan orchestrator may determine a sync-replication configuration for the one or more objects to be included in the data pool when the disaster recovery plan is performed based on the specified RPO and the specified RTO included in the data protection intent specification, and the data protection intent specification analysis. Sync-replication is the process of synchronously replicating data that is backed up data from the backup system to a destination system, such as a disaster recovery site. If the specified RTO is less than or equal to a RTO threshold duration (e.g., 5 minutes), the sync-replication configuration may require a hot standby to be created for some or all of the one or more objects (e.g., the CDP-enabled objects) to be included in the data pool when the disaster recovery plan is performed. If the specified RTO is greater than the RTO threshold duration, the sync-replication configuration may not require a hot standby to be created for some or all of the one or more objects included in the data pool when the disaster recovery plan is performed.

(36) The determined sync-replication configuration may indicate a frequency at which write IOs stored in the log maintained by the CDP service are streamed to a hot standby based on the specified RPO. If the specified RPO is less than or equal to a RPO threshold duration (e.g., 5 minutes), the sync-replication configuration may require write IOs logged by the CDP service associated with the backup system to be streamed from the backup system to the hot standby as the write IOs are received by the CDP service associated with the backup system from the source system. If the specified RPO is greater than the RPO threshold duration, the sync-replication configuration may require write IOs logged by the CDP service associated with the backup system to be streamed from the backup system to the hot standby at a frequency that allows the specified RPO to be achieved (e.g., every 30 minutes, hourly, etc.). The backup system may batch the write IOs logged by CDP service associated with the backup system and send a batch of write IOs to the disaster recovery site system hosting the hot standby.

(37) The data management plan orchestrator may determine a hot-standby configuration for the one or more objects to be included in the data pool when the disaster recovery plan is performed based on the specified RPO and the specified RTO included in the data protection intent specification, and the data protection intent specification analysis. In some embodiments, the data management plan orchestrator may determine that the specified RPO and the specified RTO may be achieved by generating at a disaster recovery site system specified in the data protection intent specification a corresponding hot standby for some or all of the one or more objects to be included in the data pool. In some embodiments, the data management plan orchestrator may determine that the specified RPO and the specified RTO may be achieved by generating on the backup system a corresponding hot standby for some or all of the one or more objects to be included in the data pool when the disaster recovery plan is performed. The data management plan orchestrator may establish this hot-standby configuration even though it may not be required by the data protection intent specification because the disaster recovery plan specified by the data protection intent specification may be unable to be satisfied or and/or reduces overall system resources needed to setup and achieve the disaster recovery plan specified by the data protection intent specification.

(38) The data management plan orchestrator may implement the sync-replication configuration and the hot-standby configuration by sending to the backup system one or more commands to replicate data associated with one or more objects determined to need a hot standby to achieve the specified RTO to a recovery system. In response to receiving the one or more commands, the backup system may provide to the recovery system one or more instructions to generate a hot standby for each of the one or more objects determined to need a hot standby and data included in a reference snapshot associated with each of the one or more objects determined to need a hot standby. In response to receiving the one or more instructions and the data, the recovery system may generate a corresponding hot standby for each of the one or more objects determined to need a hot standby and populate the corresponding hot standbys with the data. The backup system may send to the recovery system write IOs logged by the CDP service associated backup system based on the determined sync-replication configuration. In response to receiving a write IO or a group of write IOs, the recovery system may update a hot standby to which the write IO applies.

(39) Techniques to update a hot standby are disclosed herein. In some embodiments, the hot standby may be rolled back to a specified recovery time point. Using the disclosed techniques, a hot standby corresponding to a CDP-enabled object may be rolled back or forward to any point in time for which there is an associated entry stored in the log maintained by the CDP service of the backup system. A hot standby corresponding to CDP-enabled object hosted on a source system may need to be rolled back to a specified recovery time point for one or more reasons. For example, the CDP-enabled object may have become infected with malware, subject to ransomware, critical data may have been permanently deleted, etc. Other systems may restore an object from a backup. However, restoring the object from a backup may not satisfy an RTO associated with the object. Other systems may identify a restoration point that occurred before the specified moment in time, generate a standby that corresponds to the identified restoration point, and apply the write IOs included in the log from the restoration point to the specified moment in time. However, this process may take a long time to perform and may be stale (e.g., the standby is missing more recent write IOs) by the time the standby is ready to be used. In contrast, the disclosed techniques herein identify a minimal set of data chunks that need to be applied to the hot standby to roll back the hot standby to the specified moment in time in a manner that allows the RTO associated with the object to be achieved.

(40) Rolling back the hot standby that corresponds to the CDP-enabled object to a specified recovery time point may include cloning the hot standby (e.g., fast clone). The backup system may send to a recovery system hosting the hot standby one or more instructions to clone the hot standby. In response to receiving the one or more instructions, the recovery system may clone the hot standby. For the example, the hot standby clone may be a VM clone. The backup system may still be streaming write IOs associated with the CDP-enabled object to the hot standby that corresponds to the CDP-enabled object. Cloning the hot standby may pause a state of the hot standby and prevent the streamed IOs from being applied to the hot standby. The backup system may identify a log entry that corresponds to the specified recovery time point. The backup system may generate a view that corresponds to the specified recovery time point by identifying a closest reference restoration point that occurs before the specified recovery time point, cloning the view that corresponds to the closest reference restoration point, and applying to the cloned view the logged write IOs that occurred from the closest reference restoration point to the specified recovery time point.

(41) The backup system may generate a view that corresponds to a current state of the cloned hot standby by identifying a closest reference restoration point that occurs before a point in time that corresponds to the current state of cloned hot standby, cloning the view that corresponds to the closest reference restoration point, and applying to the cloned view the logged write IOs that occurred from the closest reference restoration point to point in time that corresponds to the current state of the cloned hot standby.

(42) The backup system may determine the one or more differences between the view that corresponds to the specified recovery point in time and the view that corresponds to the current state of the cloned hot standby by traversing the views to determine one or more non-shared data chunks. The backup system may determine which of the one or more non-shared data chunks are associated with the view that corresponds to the specified recovery point in time. The backup system may send the one or more determined data chunks to a recovery system hosting the hot standby (e.g., the disaster recovery site system, the backup system). In response to receiving the one or more determined data chunks, the recovery system hosting the hot standby may apply the one or more received data chunks to the cloned hot standby to roll back the cloned hot standby to the specified moment in time. Subsequently, a user associated with the entity or an application associated with the entity may be given access to the cloned hot standby.

(43) In some embodiments, the hot standby may be rolled forward to a specified recovery time point. The hot standby may be rolled forward to a specified recovery time point for one or more reasons. For example, there may be a gap in write IOs associated with a CDP-enabled object that are logged by a backup system CDP service. As a result, the hot standby corresponding to a CDP-enabled object may be in an inconsistent state with respect to the CDP-enabled object hosted on the source system because one or more missing write IOs may not be applied to the hot standby corresponding to the CDP-enabled object.

(44) A backup system may receive from a source system a stream of write IOs. Each of the write IOs is associated with a sequence number. The backup system may determine that the stream of write IOs is missing one or more write IOs based on the sequence numbers of the stream of write IOs. For example, the sequence number may be represented as [incarnation ID, change ID] and the stream of logged IOs may have a sequence of [0, 0], [0, 1], [0, 2], [0, 3], [0, 5]. In this example, an IO with a sequence number of [0, 4] is missing.

(45) A CDP service associated with the backup system may inspect a log for the stream of write IOs associated with a CDP-enabled object and determine that one or more write IOs are missing from the log (e.g., a gap) based on the corresponding sequence numbers of the write IOs. In response to determining that a gap exists, the backup system may request the source system to perform a reference snapshot (e.g., a full snapshot) for the CDP-enabled object. In response to the request, the source system may insert a snapshot marker in the stream of IOs associated with the CDP-enabled object and perform the reference snapshot. The snapshot marker may indicate that a full reference restoration point was generated for the CDP-enabled object. The reference snapshot may include all of the data associated with a CDP-enabled object (e.g., metadata and content data). In response to receiving the reference snapshot, the backup system may perform a backup of the received data and generate a view that corresponds to the reference restoration point.

(46) The backup system may send to the recovery system hosting the hot standby one or more instructions to clone the hot standby and in response to the one or more instructions, the recovery system hosting the hot standby may clone the hot standby. For example, the hot standby clone may be a VM clone. The backup system may still be streaming write IOs to the hot standby. Cloning the hot standby may pause a state of the hot standby and prevent the streamed IOs from being applied to the hot standby. A log entry that corresponds to a current state of the cloned hot standby may be identified. The backup system may generate a view that corresponds to a current state of the cloned hot standby by identifying a closest reference restoration point that occurs before the current state of cloned hot standby, cloning the view that corresponds to the closest reference restoration point, and applying to the cloned view the logged write IOs that occurred from the closest reference restoration point to the current state of the cloned hot standby.

(47) The backup system may determine the one or more differences between the view that corresponds to the reference restoration point and the view that corresponds to the current state of the hot standby by traversing the views to determine one or more non-shared data chunks. The backup system may determine which of the one or more non-shared data chunks are associated with the view that corresponds to the reference restoration point. The backup system may send the one or more determined data chunks to the recovery system hosting the hot standby. In response to receiving the one or more determined data chunks, the recovery system hosting the hot standby may apply the one or more received data chunks to the cloned hot standby to fill in the gaps. Subsequently, the cloned hot standby is in a state to which the CDP-enabled object hosted on the source system may be replicated. The backup system may continue to stream one or more write IOs included in a log from the backup system to the cloned hot standby.

(48) In some embodiments, one or more write IOs included in the stream may not be applied to the hot standby. The recovery system hosting the hot standby may include a CDP service. The CDP service may apply one or more write IOs included in the stream to the hot standby. The CDP service may determine that one or more write IOs were not applied to the hot standby based on a sequence number associated with the one or more write IOs. For example, the CDP service may determine a gap exists in the one or more write IOs based on the sequence numbers associated with the one or more write IOs. The recovery system hosting the hot standby may determine the last write IO that was applied to the hot standby before the one or more missing write IOs based on a sequence number associated with the last write IO. The recovery system hosting the hot standby may send to the backup system a command to replay the stream of write IOs from the sequence number associated with the last write IO that was applied to the hot standby. In response to receiving the command, the backup system may replay the stream of write IOs from the sequence number

associated with the last write IO that was applied to the hot standby by sending to the recovery system hosting the hot standby the one or more write IOs that occurred after the last write IO that was applied to the hot standby. In response to receiving the one or more write IOs, the CDP service of the system hosting the hot standby may apply the one or more write IOs to the hot standby.

(49) FIG. 1A is a block diagram illustrating a system for generating hot standbys for continuous data protection-enabled objects in accordance with some embodiments. In the example shown, system **100** enables user **101** to establish, via a graphical user interface **103** of user device **102**, a data protection intent for one or more objects. An object may have one or more sub-objects.

(50) User device **102** may be a computer, a desktop, a laptop, a smartphone, a tablet, or any other computing device with a graphical user interface. User device **102** is associated with user **101**. User **101** may be associated with an entity, such as an individual, an enterprise, a government, a company, an organization, etc. The entity may store a plurality of objects in datacenters **142a**, **142b**. An object may correspond to a file, a database, a VM, an application, a container, a pod, etc. Although FIG. 1 depicts the entity being associated with two datacenters, an entity may be associated with one or more datacenters.

(51) User **101** may specify, via GUI **103**, a data protection intent for the one or more objects. The data protection intent includes an indication of a data pool and one or more data management services (e.g., backup, restore, disaster recovery, replication, migration, analytics, compliance, security, etc.) that are to be applied to the data pool.

(52) A data pool is a description of one or more objects to be included in the data pool. When one or more data management services are performed, the one or more data management services are performed with respect to the one or objects specified by the data pool. The description may indirectly and/or directly identify the one or more objects. An object may be associated with one or more tags, labels, features, and/or prefixes. User **101** may specify the one or more objects to be included in the data pool when the one or more data management services are performed based on a tag, label, feature, and/or prefix associated with an object.

(53) A data management provider may establish a data management as a service (DmaaS) infrastructure in cloud environment **110** (e.g., public cloud, private cloud, hybrid cloud, etc.) provided by a cloud provider (e.g., Amazon Web Services™, Microsoft™ Azure, Google Cloud™, etc.). The DmaaS infrastructure may enable entities to remotely specify a data protection plan for one or more objects associated with an entity. The DmaaS infrastructure may include control plane **112** and one or more data planes, such as data planes **121a**, **121b**.

(54) Control plane **112** may establish one or more data planes for an entity based on a location of the entity's datacenters. For example, datacenter **142a** may be located at a first location (e.g., California) and datacenter **142b** may be located at a second location (e.g., New York). The data management provider may establish data plane **121a** for a first region (e.g., “West Coast”) that includes the first location and data plane **121b** for a second region (e.g., “East Coast”) that includes the second location. A data plane may include one or more cloud storage systems that are each comprised of one or more cloud servers (e.g., EC2 instances). A data plane may be configured to control how data associated with an entity is stored in cloud storage system **122a**, cloud storage system **122b**, and/or cloud storage **132**.

(55) Datacenters **142a**, **142b** include corresponding client side component(s) **145a**, **145b**. The one or more client side components **145a**, **145b** are configured to receive from control plane **112** one or more commands to perform one or more data management services. For example, the one or more client side components **145a**, **145b** may receive one or more commands to back up content from corresponding source systems **143a**, **143b** and/or corresponding backup systems **144a**, **144b** to cloud storage **132** (e.g., Amazon Simple Storage Service (S3)). Client side components **145a**, **145b** may receive one or more commands to replicate objects having a particular tag to datacenter **142b**, **142a**, respectively. Client side components **145a**, **145b** may receive one or more commands to cause an object having a particular tag that is hosted on source systems **143a**, **143b**, respectively, to have a particular RPO and RTO.

(56) A client side component, such as client side components **145a**, **145b**, may backup to cloud storage **132** data associated with a source system, such as source systems **143a**, **143b**, or data associated with a backup system, such as backup systems **144a**, **144b**. A cloud storage system, such as cloud storage systems **122a**, **122b**, may generate metadata for the data that is stored at cloud storage **132**. In some embodiments, a cloud storage system, such as cloud storage system **122a**, **122b**, may store the generated metadata at cloud storage **132**. In some embodiments, the generated metadata is stored in cloud storage system **122a**, **122b**. In some

embodiments, a hot standby is generated at either cloud storage system **122a**, **122b**, or cloud storage **132**. (57) Control plane **112** is configured to manage the DmaaS infrastructure that enables users, such as user **101**, to remotely access and manage one or more objects associated with an entity. Control plane **112** includes a user interface (UI) service virtualization module **113** that enables user device **102** to communicate with the DmaaS infrastructure. UI service virtualization module **113** may receive from device **102** via GUI **103** a data protection intent specification. The data protection intent specification may include a disaster recovery plan for the one or more objects to be included in the data pool when the disaster recovery plan is performed. The disaster recovery plan may indicate a primary site associated with each of the one or more objects, a disaster recovery site associated with each of the one or more objects, a specified RPO for the one or more objects, and/or a specified RTO for the one or more objects. UI service virtualization module **113** may provide the data protection intent specification to data management plan orchestrator **114**.

(58) Data management plan orchestrator **114** may analyze the data protection intent specification to determine configurations for one or more data storage services needed to achieve the data protection intent specification. The one or more data storage services may include protection jobs, sync-replication, and/or hot-standby. The data protection intent specification analysis may identify one or more types of objects specified to be included in the data pool when the disaster recovery plan is performed. For example, data management plan orchestrator **114** may identify, based on tags, labels, and/or features associated with a particular object (e.g., file size, creation date, ownership, last modification date, etc.), whether the particular object is a high IO object (e.g., an object having a number of IOs/min greater than or equal to an IO threshold), whether the particular object is a low IO object (e.g., an object having a number of IOs/min less than or equal to the IO threshold), performance requirements associated with the particular object, etc. In some embodiments, data management plan orchestrator **114** implements a machine learning model (e.g., supervised, unsupervised, semi-supervised, regression, etc.) to identify based on tags, labels, and/or features associated with a particular object (e.g., file size, creation date, ownership, last modification date, etc.), whether the particular object is a high input/output(IO) object (e.g., an object having a number of IOs/min greater than or equal to an IO threshold), whether the particular object is a low IO object (e.g., an object having a number of IOs/min less than or equal to the IO threshold), performance requirements associated with the particular object, etc.

(59) Data management plan orchestrator **114** may determine protection jobs for the one or more objects to be included in the data pool when the disaster recovery plan is performed based on a specified RPO included in the data protection intent specification and the data protection intent specification analysis. In some embodiments, data management plan orchestrator **114** determines that a scheduled backup (e.g., every hour, every six hours, daily, etc.) is sufficient to achieve the specified RPO for an object. In some embodiments, data management plan orchestrator **114** determines that CDP needs to be enabled to achieve the specified RPO for an object.

(60) Data management plan orchestrator **114** may automatically enable CDP for one or more objects to be included in the data pool when the disaster recovery plan is performed that are determined to need CDP to achieve a specified RPO by providing to a backup system associated with the source system (e.g., backup systems **144a**, **144b** are associated with source systems **143a**, **143b**, respectively) one or more commands to enable CDP for the one or more objects determined to need CDP to achieve the specified RPO. In response to receiving the one or more commands, a backup system, such as backup systems **144a**, **144b**, may provide to a source system, such as source systems **143a**, **143b**, a request to perform a corresponding reference snapshot and a corresponding write filter for each of the one or more objects determined to need CDP to achieve the specified RPO and enable on the backup system a CDP service to handle write IOs. In some embodiments, source systems **143a**, **143b** include a corresponding write filter for each of the one or more objects and the one or more commands enable the corresponding write filters.

(61) In response to receiving the request to perform a corresponding reference snapshot, the source system may perform a full backup snapshot and include all of the object data associated with the one or more objects determined to need CDP to achieve the specified RPO. The backup system may ingest the data included in the reference snapshot and generate a tree data structure for the one or more objects determined to need CDP to achieve the specified RPO. In some embodiments, the tree data structure includes a snapshot tree for an object and one or more metadata structures that corresponding to one or more object files associated with the object. In some embodiments, the tree data structure is a snapshot tree for all of the

objects hosted on the source system and a corresponding metadata structure for each of the one or more objects hosted by the source system that are to be included in the data pool when the disaster recovery plan is performed.

(62) In response to receiving a write filter, the source system may attach the corresponding write filter to each of the one or more objects determined to need CDP to achieve the specified RPO. A write filter may be configured to intercept all IOs of an associated object, filter the read IOs, and provide a stream of the intercepted write IOs to the backup system. The write filter may be configured to attach a sequence number to each intercepted write IO. In some embodiments, the source system hosting the one or more objects is capable of providing CDP for the one or more objects because the source system is capable of providing its own write filters. The backup system may provide to the source system a command to enable CDP for some or all of the one or more objects. In response to receiving the command, the source system may attach a corresponding write filter to each of the one or more objects determined to need CDP to achieve the specified RPO and corresponding write filters may stream the write IOs to a backup system.

(63) In some embodiments, a backup system, such as backup system **144a**, **144b**, is configured to provide an IO filter daemon to a source system, such as source system **143a**, **143b**. The IO filter daemon of the source system may be configured to receive the intercepted write IOs from each of the write filters of the source system and to provide the intercepted write IOs from each of the write filters of the source system to the backup system.

(64) A backup system, such as backup systems **144a**, **144b**, may include one or more CDP services that are configured to receive one or more intercepted write IOs from the IO filter daemon of a source system, such as source systems **143a**, **143b**. Each cluster node of the backup system may have a corresponding CDP service. The IO filter daemon may select one of the one or more CDP services and provide the intercepted write IOs associated with a particular object to the selected CDP service. A CDP service may maintain a log for the one or more intercepted write IOs received from the IO filter daemon of the source system. Each entry of the log may store the metadata associated with an intercepted write IO. For example, the metadata associated with an intercepted write IO may include at least one of the following: sequence number, timestamp, IO offset, IO length, IO data bytes, VM UUID, VMDK full path on a host at which the VM is running, etc.

(65) Data management plan orchestrator **114** may determine a sync-replication configuration for the one or more objects to be included in the data pool when the disaster recovery plan is performed based on the specified RPO and the specified RTO included in the data protection intent specification, and the data protection intent specification analysis. If the specified RTO is less than or equal to an RTO threshold duration (e.g., 5 minutes), the sync-replication configuration may require a hot standby to be created for some or all of the one or more objects to be included in the data pool when the disaster recovery plan is performed. If the specified RTO is greater than the RTO threshold duration, the sync-replication configuration may not require a hot standby to be created for some or all of the one or more objects included in the data pool when the disaster recovery plan is performed.

(66) The determined sync-replication configuration may indicate a frequency at which write IOs stored in the log maintained by the CDP service of a backup system, such as backup systems **144a**, **144b**, are streamed to a hot standby based on the specified RPO. If the specified RPO is less than or equal to a RPO threshold duration (e.g., 5 minutes), the sync-replication configuration may require write IOs logged by the CDP service associated with the backup system to be streamed from the backup system to the hot standby as the write IOs are received by the CDP service associated with the backup system from the source system. If the specified RPO is greater than the RPO threshold duration, the sync-replication configuration may require write IOs logged by the CDP service associated with the backup system to be streamed from the backup system to the hot standby at a frequency that allows the specified RPO to be achieved (e.g., every 30 minutes, hourly, etc.). The backup system may batch the write IOs logged by CDP service associated with the backup system and send a batch of write IOs to the disaster recovery site system hosting the hot standby.

(67) Data management plan orchestrator **114** may determine a hot-standby configuration for the one or more objects to be included in the data pool when the disaster recovery plan is performed based on the specified RPO and the specified RTO included in the data protection intent specification, and the data protection intent specification analysis. In some embodiments, data management plan orchestrator **114** may determine that the specified RPO and the specified RTO may be achieved by generating at a disaster

recovery site system specified in the data protection intent specification a corresponding hot standby for each of the one or more objects to be included in the data pool. For example, source system **143b** or backup system **144b** may be a disaster recovery site system for datacenter **142a** and source system **143a** or backup system **144a** may be disaster recovery site system for datacenter **142b**. In some embodiments, data management plan orchestrator **114** may determine that the specified RPO and the specified RTO may be achieved by generating on at a different disaster recovery site a corresponding hot standby for each of the one or more objects to be included in the data pool when the disaster recovery plan is performed. For example, data management plan orchestrator **114** may determine to generate for an object hosted on source system **143a** a hot standby at a different disaster recovery site instead of determining to generate for the object hosted on source system **143a** a hot standby on either source system **143b** or backup system **144b**. Data management plan orchestrator **114** may establish this hot-standby configuration even though it does not follow the data protection intent specification because the disaster recovery plan specified by the data protection intent specification may be unable to be satisfied or and/or reduces overall system resources needed to setup and achieve the disaster recovery plan specified by the data protection intent specification.

(68) Data management plan orchestrator **114** may validate the data protection intent specification that includes the disaster recovery plan. Data management plan orchestrator **114** may determine whether the specified data protection intent is able to be achieved. In the event the specified data protection intent is unable to be achieved, data management plan orchestrator **114** may notify, via GUI **103** of device **102**, user **101** that the data protection intent is unable to be achieved. In the event the intent of user **101** is able to be achieved, data management plan orchestrator **114** may orchestrate the datacenter components of system **100** that are needed to achieve the intent (e.g., client side components **145a**, **145b**, source systems **143a**, **143b**, backup systems **144a**, **144b**, cloud storage systems **122a**, **122b**, and/or cloud storage **132**). For example, after validation, data management plan orchestrator may send one or more commands to one or more data storage devices needed to implement the determined configurations of the one or more data storage services to achieve the data protection intent.

(69) Data management plan orchestrator **114** may implement the sync-replication configuration and the hot-standby configuration by sending to a backup system, such as backup systems **144a**, **144b**, one or more commands to replicate data associated with one or more objects determined need a hot standby to achieve the specified RTO to a recovery system. In response to receiving the one or more commands, the backup system may provide to the recovery system one or more instructions to generate a hot standby for each of the one or more objects determined to need a hot standby and data included in a reference snapshot associated with each of the one or more objects determined to need a hot standby. In response to receiving the one or more instructions and the data, the recovery system may generate a corresponding hot standby for each of the one or more objects determined to need a hot standby and populate the corresponding hot standbys with the data. In some embodiments, a VM hot standby is generated for each of the one or more objects determined to need a hot standby. The backup system may send to the recovery system write IOs logged by the CDP service associated backup system based on the determined sync-replication configuration. In response to receiving a write IO, the recovery system may update a hot standby to which the write IO applies.

(70) In some embodiments, the data that is being backed up to a backup system is also being replicated, in parallel, to a recovery system. For example, data that is being backed up from source system **143a** to backup system **144a** is replicated, in parallel, to source system **143b**. Source system **143b** may include a hot standby that is configured to apply the data so that the hot standby mirrors a CDP-enabled object hosted on the source system.

(71) FIG. 1B is a block diagram of a system for providing continuous data protection in accordance with some embodiments. In the example shown, system **150** is comprised of source system **152** that is coupled to backup system **162** via connection **160**. Connection **160** may be a wired or wireless connection. Connection **160** may be a LAN, WAN, intranet, the Internet, and/or a combination thereof.

(72) Source system **152** may be comprised of one or more objects **153**, a backup agent **154**, one or more write filters **105**, and a filter daemon **106**. Source system **152** may be configured to host one or more objects **153**. Source system **152** may include one or more storage volumes (not shown) that are configured to store file system data associated with source system **152**. The file system data associated with source system **152** includes the data associated with the one or more objects **153**.

(73) Backup agent **154** may be configured to cause source system **152** to perform a reference snapshot (e.g.,

a full snapshot). A full snapshot may include all of the file system data of source system **152** at a particular moment in time. In some embodiments, a full snapshot includes all of the object data associated with some or all of the one or more objects **153** at a particular moment in time. In some embodiments, backup agent **154** is running on source system **152**. In some embodiments, backup agent **154** is running in one of the one or more objects **153**. In some embodiments, a backup agent **154** is running on source system **152** and a separate backup agent **154** is running in one of the one or more objects **153**.

(74) In some embodiments, an object includes a backup function and is configured to perform a reference snapshot on its own without backup agent **154**. In some embodiments, source system **152** includes a backup function and is configured to perform a reference snapshot on its own without backup agent **154**. In some embodiments, backup system **162** may provide instructions to source system **152**, causing source system **152** to execute backup functions without backup agent **154**.

(75) Backup system **162** may receive the file system data included in a reference snapshot and generate a view of the file system data. In some embodiments, a view of the file system data corresponding to the reference snapshot is generated for source system **152**. In some embodiments, a view of the object data corresponding to the reference snapshot is generated for some or all of the one or more objects **153**.

(76) File system manager **167** may be configured to organize in a tree data structure the file system data and/or object data included in the reference snapshot. An example of the tree data structure is a snapshot tree, which may be based on a B+ tree structure (or other type of tree structure in other embodiments). In the event the reference snapshot corresponds to all of the file system data of source system **152**, the view corresponding to the reference snapshot may be comprised of a snapshot tree and one or more object metadata structures. The snapshot tree may be configured to organize the file system data associated with source system **152**. An object metadata structure may be configured to store the metadata associated with one of the one or more objects **153**. The metadata associated with one of the one or more objects **153** may enable data chunks associated with the one or more objects **153** that were backed up to backup system **162** to be located. Each of the one or more objects **153** may have a corresponding metadata structure.

(77) In the event the reference snapshot corresponds to all of the object data of one of the one or more objects **153**, the view corresponding to the reference snapshot may be comprised of a snapshot tree and one or more object file metadata structures. The snapshot tree may be configured to store the metadata associated with one of the one or more objects **153**. An object file metadata structure may be configured to store the metadata associated with an object file included in the object that enables the data chunks associated with the object to be located.

(78) The tree data structure may be used to capture different reference restoration points. The tree data structure allows a chain of reference restoration points to be linked together by allowing a node of a later version of a snapshot tree associated with an intermediate reference restoration point to reference a node of a previous version of a snapshot tree. The previous version of the snapshot tree may correspond to an earlier intermediate reference restoration point or an earlier full reference restoration point. For example, a root node or an intermediate node of a snapshot tree corresponding to the later reference restoration point may reference an intermediate node or leaf node of a snapshot tree corresponding to the earlier reference restoration point.

(79) A snapshot tree is a representation of a fully hydrated restoration point because it provides a complete view of a source system or an object at a particular moment in time. A fully hydrated restoration point is a restoration point that is ready for use without having to reconstruct a plurality of backups to use it. Instead of reconstructing a restoration point by starting with a full backup and applying one or more data changes associated with one or more incremental backups to the data associated with the full backup, backup system **162** maintains fully hydrated restoration points. Any file associated with source system **152** or an object at a particular time and the file's contents, for which there is an associated reference restoration point, may be determined from the snapshot tree, regardless if the associated reference restoration was a full reference restoration point or an intermediate reference restoration point.

(80) Backup system **162** may generate a fully hydrated restoration point that corresponds to any point in time for which there is an associated log entry. Backup system **162** may identify a closest reference restoration point that occurred before the point in time and apply write IOs included in the log from the closest reference restoration point to the point in time.

(81) A snapshot tree may include a root node, one or more levels of one or more intermediate nodes associated with the root node, and one or more leaf nodes associated with an intermediate node of the

lowest intermediate level. The root node of a snapshot tree may include one or more pointers to one or more intermediate nodes. Each intermediate node may include one or more pointers to other nodes (e.g., a lower intermediate node or a leaf node). A leaf node may store file system metadata, data associated with a file that is less than a limit size, an identifier of a data brick, a pointer to a metadata structure (e.g., object metadata structure or an object file metadata structure), a pointer to a data chunk stored on the storage cluster, etc.

(82) A metadata structure (e.g., object file metadata structure or object metadata structure) may include a root node, one or more levels of one or more intermediate nodes associated with the root node, and one or more leaf nodes associated with an intermediate node of the lowest intermediate level. The tree data structure associated with a metadata structure allows a chain of metadata structures corresponding to different versions of an object or an object file to be linked together by allowing a node of a later version of a metadata structure to reference a node of a previous version of a metadata structure.

(83) A metadata structure may be associated with a plurality of chunk files. A chunk file may be comprised of a plurality of data chunks. Metadata store **164** may store one or more metadata structures. In some embodiments, a metadata structure corresponds to an object. In other embodiments, a metadata structure corresponds to an object file. In other embodiments, a metadata structure corresponds to a portion of an object.

(84) A leaf node of a metadata structure may store a value, such as an identifier of a data brick associated with one or more data chunks. The location of the one or more data chunks associated with a data brick may be identified using one or more data structures (e.g., list, table, etc.) stored in metadata store **164**. A first data structure (e.g., chunk metadata table) may store information that associates a brick identifier with one or more chunk identifiers and one or more chunk file identifiers. A second data structure (e.g., chunk file metadata table) may associate a chunk file identifier with a chunk file storing a plurality of data chunks. In some embodiments, the first data structure and the second data structure are combined as a single data structure. The one or more data chunks associated with a data brick may be located based on the chunk metadata table and the chunk file metadata table. For example, a first data brick having a first brick identifier may be associated with a first chunk identifier (e.g., SHA-1 hash value). The first chunk identifier may be used in conjunction with the chunk metadata table to identify a chunk file identifier. A chunk file having the identified chunk file identifier is comprised of a plurality of data chunks. The chunk file metadata table may be used to identify a location of the plurality of data chunks. The chunk file metadata table may include offset information of the plurality of data chunks within a chunk file.

(85) One or more write filters **155** may be attached to the one or more objects **153** prior to a reference snapshot (e.g., full snapshot) of source system **152** being initiated to provide CDP for the object to which a write filter is attached. The write filter may be attached at a disk layer of the object (e.g., disk write filter), a file system layer of the object (e.g., file system write filter), or a virtualized layer of the object (e.g., virtualized write filter). An object may be associated with one or more disks. A corresponding write filter may be attached to each of the one or more disks. A disk may be associated with one or more write filters.

(86) Each object of the one or more objects **153** may be associated with a corresponding write filter. A write filter may be configured to intercept all IOs of an associated object, filter out the read IOs, and provide a stream of the intercepted write IOs. A write filter may be configured to attach a sequence number to each intercepted write IO. The sequence number may include an incarnation ID and a change ID. The incarnation ID may indicate an incarnation of the write filter and may be incremented each time the write filter is restarted after a crash (e.g., object crash, filter crash). In some embodiments, the incarnation ID is not incremented after a clean restart of an object. The change ID may be incremented for each intercepted write IO. For example, the sequence number may be represented as [incarnation ID, change ID]. The sequence number may be [0, 0], [0, 1] . . . [0, n]. In the event the write filter goes offline after a crash and comes back online, the incarnation ID may be incremented such that the sequence numbers attached to the intercepted write IOs may be [1, 0], [1, 1] . . . [1, n].

(87) The one or more write filters **155** may be configured to provide their corresponding intercepted write IOs to filter daemon **156**. Filter daemon **156** may be configured to collect the intercepted write IOs from the plurality of write filters **155** and stream the intercepted write IOs to backup system **162**. In some embodiments, filter daemon **156** streams the intercepted write IOs as they are received. In some embodiments, filter daemon **156** is configured to stream batches of intercepted write IOs. Filter daemon **156** may be configured to cache the received intercepted write IOs until they are provided to backup system

162.

(88) Filter daemon **156** may be configured to provide the intercepted write IOs associated with a particular object of the one or more objects **153** to a particular cluster node of cluster nodes **161, 163, 165**. The number of available cluster nodes may increase (e.g., a node is added to backup system **162**) or decrease (e.g., a node is offline). A particular cluster node may be assigned to a particular object. Filter daemon **106** may determine an object_ID associated with an object and stream the intercepted write IOs associated with the object_ID to a cluster node assigned to an object having the object_ID. The particular cluster node may be selected from the set of available cluster nodes based on one or more factors, such as availability of a CDP service associated with a cluster node.

(89) Filter daemon **156** may be configured to stream the intercepted write IOs associated with a plurality of objects to corresponding cluster nodes of cluster nodes **161, 163, 165**. For example, filter daemon **156** may be configured to stream the intercepted write IOs associated with a first object to cluster node **161**, the intercepted write IOs associated with a second object to cluster node **163**, . . . , and the intercepted write IOs associated with an nth object to cluster node **165**. The intercepted write IOs associated with a plurality of objects may be streamed in parallel. In some embodiments, the intercepted write IOs associated with a plurality of objects are streamed sequentially. In some embodiments, the intercepted write IOs associated with an object are streamed to a plurality of cluster nodes.

(90) In some embodiments, the intercepted write IOs associated with a particular object are grouped into a plurality of sets, e.g., a first set of intercepted write IOs, a second set of intercepted write IOs, . . . , and an nth set of intercepted write IOs. The different sets of IOs may be provided to different nodes of cluster nodes **161, 163, 165**. For example, the first set of intercepted write IOs associated with the particular object may be provided to cluster node **161**, the second set of intercepted write IOs associated with the particular object may be provided to cluster node **163**, . . . , and the nth set of intercepted write IOs associated with the particular object may be provided to cluster node **165**. The intercepted write IOs associated with the plurality of sets may be streamed in parallel. In some embodiments, the intercepted write IOs associated with the plurality of sets are streamed sequentially.

(91) Backup agent **154** may cause source system **152** to perform an incremental snapshot. An incremental snapshot usually includes the data that was not previously backed up. However, the one or more write filters **105** and filter daemon **106** are configured to provide CDP for the one or more objects **153**. Thus, a stream of data associated with the one or more objects **153** is continuously being provided from source system **152** to backup system **162**. Instead of providing a command to perform an incremental snapshot, backup agent **154** may provide to an object a command to generate an intermediate reference restoration point. Instead of performing the incremental backup snapshot, a write filter may intercept the command as it is being provided to the object and attach a snapshot marker to the IO. The snapshot marker indicates a time at which an intermediate reference restoration point associated with the object was requested. The snapshot marker provides an intermediate reference restoration point, which may be used to identify a version of an object to be used to restore the object to a particular point in time.

(92) In some embodiments, an application associated with an object may perform a snapshot of the object. The snapshot has an associated snapshot ID that is not identified by backup system **162**. A snapshot marker has a particular snapshot ID. A snapshot marker corresponding to a snapshot initiated by backup system **162** may have an associated ID that is known to backup system **162**. A snapshot marker corresponding to a snapshot initiated by an application associated with the object may have an associated ID that is unknown to backup system **162**. In the event an application associated with an object performs a snapshot of the object, the snapshot may not be used as a reference restoration point because the backup system is unaware of a snapshot having the ID that is associated with the application.

(93) Backup system **162** is comprised of a storage cluster that includes a plurality of cluster nodes **161, 163, 165**. Although three cluster nodes are shown, backup system **162** may be comprised of n nodes. In some embodiments, a cluster node includes a processor and memory, and is coupled to a separate storage device. The separate storage device may include one or more storage devices (e.g., flash storage devices). A storage device may be segmented into a plurality of partitions. In some embodiments, one of the partitions is configured to store all of the backup system **162** generated metadata (e.g., snapshot trees, object metadata structures) included in metadata store **164**. In some embodiments, a partition is configured to store some of the metadata store **164** and the backup system **162** generated metadata included in metadata store **164** is stored in a plurality of partitions. A node may be allocated one of the partitions that stores some of metadata

store **164**. Each of the cluster nodes **161**, **163**, **165** may be allocated one or more of the partitions. The one or more partitions allocated to a node may be configured to store data associated with some or all of the plurality of objects that were backed up to backup system **162**. For example, the separate storage device may be segmented into 10 partitions and backup system **162** may include 10 cluster nodes. A cluster node of the 10 cluster nodes may be allocated one of the 10 partitions.

(94) In some embodiments, a cluster node includes a processor, memory, and a storage device. The storage device may be configured to store metadata associated with a plurality of objects that were backed up to backup system **162**, such as part or all of metadata store **164**. For example, the metadata associated with the plurality of objects may include backup system **162** generated metadata (e.g., metadata that enables the data associated with the object to be located). The cluster node may be coupled to a separate storage device. The separate storage device may include one or more storage devices. A storage device may be segmented into a plurality of partitions. Each of the cluster nodes **161**, **163**, **165** may be allocated one or more of the partitions. The one or more partitions allocated to a node may be configured to store data associated with some or all of the plurality of objects that were backed up to backup system **162**. For example, the separate storage device may be segmented into 10 partitions and backup system **162** may include 10 cluster nodes. A cluster node of the 10 cluster nodes may be allocated one of the 10 partitions.

(95) In some embodiments, a cluster node includes a processor, memory, and one or more storage devices. A first storage device may be configured to store metadata associated with a plurality of objects that were backed up to backup system **162**, such as part or all of metadata store **164**. For example, the metadata associated with the plurality of objects may include backup system **162** generated metadata. A second storage device may be configured to store data associated with the plurality of objects that were backed up to backup system **162**. The plurality of storage devices may include one or more solid state drives, one or more hard disk drives, or a combination thereof.

(96) Backup system **162** may be a cloud instantiation of a backup system. A configuration of cloud instantiation of backup system **162** may be a virtual replica of backup system **162**. For example, backup system **162** may be comprised of three cluster nodes, each cluster node with a storage capacity of 10 TB. A cloud instantiation of backup system **162** may be comprised of three virtual nodes, each virtual node with a storage capacity of 10 TB. In other embodiments, a cloud instantiation of backup system **162** may have more storage capacity than an on-premises instantiation of backup system **162**. In other embodiments, a cloud instantiation of backup system **162** may have less storage capacity than an on-premises instantiation of backup system **162**.

(97) Backup system **162** may include CDP Service **166**. In some embodiments, cluster nodes **161**, **163**, **165** each have a corresponding CDP Service. In other embodiments, backup system **162** has a single CDP service that is running on one of the cluster nodes **161**, **163**, **165**. CDP Service **166** may be configured to receive one or more intercepted write IOs from filter daemon **156**. Filter daemon **156** may assign a particular CDP Service **166** for a particular object. CDP Service **166** may have an associated cache. In response to receiving the one or more intercepted write IOs, CDP Service **166** may be configured to store in the associated cache, the one or more received intercepted write IOs in a log (e.g., write-ahead-log). CDP Service **166** may be configured to store a corresponding log per object. In some embodiments, CDP Service **166** is configured to store a log for a plurality of objects.

(98) A log may be comprised of one or more entries. An entry may associate an intercepted write IO stored by CDP Service **166** with a sequence number. In some embodiments, an entry associates the intercepted write IO stored by CDP Service **166** with a sequence number and a snapshot marker. File system manager **167** may determine that the log includes an entry with a snapshot marker. In response to determining that the log includes an entry with a snapshot marker, file system manager **167** may generate a reference restoration point (e.g., a view corresponding to the reference restoration point) by updating a tree data structure corresponding to the object to be restored. In some embodiments, a snapshot tree corresponding to an object is cloned and the one or more entries included in the log are applied to the cloned snapshot tree. In some embodiments, an object metadata structure is cloned and the one or more entries included in the log are applied to the cloned object metadata structure. Each entry that includes a snapshot marker corresponds to a fully hydrated reference restoration point.

(99) In some embodiments, file system manager **167** may detect a gap in a sequence of IOs associated with an object. For example, file system manager **167** may detect that an incarnation ID associated with a write filter attached to an object has incremented. File system manager **167** may detect that a change ID has been

skipped. In response to detecting a gap in the sequence of IOs associated with an object, file system manager **167** may provide to backup agent **154** a notification of the gap. In response to receiving the notification, backup agent **154** may cause a reference snapshot of the affected object to be performed. In some embodiments, in response to receiving the notification, backup agent **154** may cause a reference snapshot of source system **152**. This ensures that the data associated with an object that is backed up to backup system **162** is not corrupt because the object data is missing data.

(100) FIG. **1C** is a block diagram illustrating an example of a source system in accordance with some embodiments. In the example shown, source system **170** is comprised of a first object **171a**, a second object **171b**, . . . , and an nth object **171n**. Each object may be a different VM. Although FIG. **1C** depicts an object as a VM, source system **170** may host different objects, such as a database, an application, a container, a pod, etc. Each object is associated with a corresponding write filter. For example, object **171a** is associated with write filter **172a**, object **171b** is associated with write filter **172b**, and object **171n** is associated with write filter **172n**. A write filter may be coupled (e.g., attached) to a corresponding object prior to an initial reference snapshot (e.g., full snapshot) of the object being initiated. A write filter may be configured to intercept IOs associated with the object to which the write filter is coupled. A write filter may include a sequence number along with the intercepted write IO. In some embodiments, a write filter includes a snapshot marker in between the intercepted IOs in a stream of intercepted IOs. A write filter may be configured to provide an intercepted write IO to filter daemon **174**.

(101) Source system **170** may include a host **173** (e.g., hypervisor) that includes filter daemon **174**. Filter daemon **174** may be configured to collect the intercepted write IOs and/or snapshot marker(s) provided by write filters **172a**, **172b**, **172n**. Filter daemon **174** may be configured to stream to a backup system the intercepted write IOs and/or snapshot marker(s) as they are received. In other embodiments, filter daemon **174** is configured to stream to the backup system batches of intercepted write IOs and/or snapshot marker(s).

(102) In some embodiments, filter daemon **174** streams the intercepted write IOs associated with a particular object to a specific cluster node of the backup system. In some embodiments, filter daemon **174** streams the intercepted write IOs associated with a particular object to a plurality of cluster nodes of the backup system.

(103) FIG. **1D** is a block diagram illustrating an example of a backup system in accordance with some embodiments. Backup system **180** includes a first cluster node **175**, a second cluster node **185**, . . . , and an nth cluster node **195**. Each cluster node has a corresponding processor, a corresponding CDP service, and a plurality of storage tiers. A first storage tier may correspond to a storage with faster access times than a second storage tier. In the example shown, a first storage tier corresponds to a solid state drive and a second storage tier corresponds to a hard disk drive.

(104) Cluster node **175** includes processor **176**, solid state drive **177**, hard disk drive(s) **178**, and CDP service **179**. Cluster node **185** includes processor **186**, solid state drive **187**, hard disk drive(s) **188**, and CDP service **189**. Cluster node **195** includes processor **196**, solid state drive **197**, hard disk drive(s) **198**, and CDP service **199**. In some embodiments, a cluster node includes a plurality of different CDP services. A CDP service may be an application, a module, a plug-in, a VM, a container, etc.

(105) Backup system **180** may receive a reference snapshot (e.g., full snapshot) from a source system. The data included in the reference snapshot may be stored in one of the cluster nodes **175**, **185**, **195**. In some embodiments, the data included in the reference snapshot is stored across the cluster nodes **175**, **185**, **195**. The data included in the reference snapshot may be stored in a first storage tier, a second storage tier, or a combination of a first storage tier and a second storage tier.

(106) Backup system **180** may include a file system manager that is configured to organize the data included in the reference snapshot. In some embodiments, a file system manager is running on one of the cluster nodes **175**, **185**, **195**. In some embodiments, each cluster node has its own corresponding file system manager. The file system manager may organize the data included in the reference snapshot using a tree data structure. The tree data structure may be configured to store the metadata associated with a source system. In some embodiments, the tree data structure is configured to store the metadata associated with an object. The tree data structure and the metadata may be stored in a metadata store. The metadata store may be distributed across the cluster nodes **175**, **185**, **195**. Solid state drives **177**, **187**, **197** may each store a portion of the metadata store.

(107) Each CDP Service may have a corresponding log. CDP Services **179**, **189**, **199** may receive a

corresponding stream of IOs from a filter daemon, such as filter daemon 174, and store a record of the IO in its corresponding log. The data associated with an IO may be stored in a storage associated with the CDP Service (e.g., a solid state drive, a hard disk drive, a cache, etc.). A log may include one or more entries. In some embodiments, the log stores one or more entries for a particular object. In some embodiments, the log stores entries for a plurality of objects. In some embodiments, a CDP service has a plurality of logs, each log corresponding to a particular object.

(108) An entry of the log may associate a particular IO with a timestamp and a sequence number. In some embodiments, the entry of the log associates a particular IO with a timestamp, a sequence number, and a snapshot marker. The snapshot marker indicates that a reference restoration point is to be generated. Each entry of the log may store the metadata associated with an intercepted write IO. For example, the metadata associated with an intercepted write IO may include at least one of the following: sequence number, timestamp, IO offset, IO length, IO data bytes, VM universally unique identifier (UUID), VM disk (VMDK) full path on a host at which the VM is running, etc.

(109) FIG. 2 is an example of a continuous data protection timeline in accordance with some embodiments. In the example shown, continuous data protection timeline 200 illustrates reference restoration points that are generated at 12:00 μ m, 12:15 μ m, 12:30 μ m, 12:45 μ m, and 1:00 μ m. Reference restoration points may be generated based on a backup policy. In some embodiments, the backup policy indicates that reference restoration points are to be generated periodically (e.g., every 15 minutes, every 30 minutes, every 60 minutes, etc.). In some embodiments, the backup policy indicates that reference restoration points are to be generated in response to a user command. In some embodiments, the backup policy indicates that reference restoration points are to be generated after a threshold number of data changes have been logged. In some embodiments, the backup policy indicates that a reference restoration point is to be generated in the event a gap is detected in the sequenced of logged IOs.

(110) In the example shown, at 12:00 μ m, a full reference restoration point is generated. A full reference restoration point may correspond to a full backup snapshot. A full backup snapshot may include all of the file system data associated with a source system comprised of one or more objects. File system data may be ingested and stored by a backup system. The backup system may include a file system manager that is configured to organize the file system data in a tree data structure. The tree data structure may be comprised of a snapshot tree and one or more metadata structures. A metadata structure may be an object metadata structure that corresponds to one of the one or more objects hosted on the source system.

(111) In some embodiments, the full backup snapshot includes all of the object data associated with an object. Object data may be ingested and stored by a backup system. The backup system may include a file system manager that is configured to organize the object data in a tree data structure. The tree data structure may be comprised of a snapshot tree and one or more metadata structures. A metadata structure may be an object file metadata structure that corresponds to one of the files associated with the object.

(112) Between 12:00 μ m and 12:15 μ m, a write filter attached to an object hosted on a source system may be configured to stream intercepted write IOs to a CDP service of the backup system. The intercepted write IOs are associated with one or more objects. In response to receiving the intercepted write IOs, the CDP service may log the intercepted write IOs.

(113) At 12:15 μ m, an intermediate reference restoration point is generated. The intermediate reference restoration point may be generated by cloning a tree data structure associated with a previous reference restoration point and applying the logged IOs to the cloned tree data structure. In this example, the previous application consistent recovery point corresponds to a full reference restoration point. The intermediate reference restoration point may be generated by cloning a root node of a tree data structure associated with the previous reference restoration point. In some embodiments, the root node of a snapshot tree associated with the previous reference restoration point is cloned. In some embodiments, the root node of a metadata structure associated with the previous reference restoration point (object metadata structure or object file metadata structure) is cloned. The cloned tree data structure may be updated, as described herein, based on the logged IOs. The logged IOs that are applied to the tree data structure associated with the previous reference restoration point are the IOs that are logged between 12:00 μ m and 12:15 μ m.

(114) Between 12:15 μ m and 12:30 μ m, a write filter attached to the object hosted on the source system may be configured to stream intercepted write IOs to a CDP service of the backup system. The intercepted write IOs are associated with an object. In response to receiving the intercepted write IOs, the CDP service may log the intercepted write IOs.

(115) At 12:30 pm, an intermediate reference restoration point is generated. The intermediate reference restoration point may be generated by cloning a tree data structure associated with a previous reference restoration point (e.g., 12:15 pm) and applying the logged IOs to the cloned tree data structure. In this example, the previous reference restoration point corresponds to an intermediate reference restoration point. The intermediate reference restoration point may be generated by cloning a root node of a tree data structure associated with the previous reference restoration point. In some embodiments, the root node of a snapshot tree associated with the previous reference restoration point is cloned. In some embodiments, the root node of a metadata structure associated with the previous reference restoration point (object metadata structure or object file metadata structure) is cloned. The cloned tree data structure may be updated, as described herein, based on the logged IOs. The logged IOs that are applied to the tree data structure associated with the previous application consistent recovery point are the IOs that are logged between 12:15 pm and 12:30 pm.

(116) Between 12:30 pm and 12:45 pm, intercepted write IOs are logged by a CDP service and at 12:45 pm, an intermediate reference restoration point is generated. The intermediate reference restoration point may be generated in the same manner as described with respect to 12:30 pm.

(117) Between 12:45 pm and 1:00 pm, intercepted write IOs are logged by a CDP service. A CDP service may inspect the log and determine that a gap exists in the sequence numbers associated with the entries that were logged between 12:45 pm and 1:00 pm. In response to detecting a gap, the CDP service may send to a backup agent of the source system a command to generate a full reference restoration point. In response to receiving the command, the backup agent may cause the source system to perform a reference snapshot.

(118) FIG. 3A is a block diagram illustrating an embodiment of a tree data structure. A tree data structure may be used to represent the file system data that is stored on a backup system, such as backup system **162**. In some embodiments, a tree data structure is used to represent the object data associated with an object that is stored on the backup system. A backup system may receive from a source system data associated with a plurality of IOs and store the data. A file system manager, such as file system manager **167**, may organize the received data in a tree data structure, such as tree data structure **300**. In the example shown, tree data structure **300** corresponds to a snapshot tree. In some embodiments, tree data structure **300** is configured to store the file system metadata associated with a source system. In some embodiments, tree data structure **300** is configured to store the object metadata associated with an object. Tree data structure **300** may correspond to a reference snapshot, e.g., a full backup snapshot. Tree data structure **300** may be a view that corresponds to a reference restoration point.

(119) Tree data structure **300** is comprised of a snapshot tree that includes a root node **302**, intermediate nodes **312**, **314**, and leaf nodes **322**, **324**, **326**, **328**, and **330**. Although tree data structure **300** includes one intermediate level between root node **302** and leaf nodes **322**, **324**, **326**, **328**, **330**, any number of intermediate levels may be implemented. Tree data structure **300** may correspond to a snapshot of file system data or object data at a particular point in time *t*. The snapshot tree depicted in FIG. 3A may correspond to a reference restoration point (e.g., full reference restoration point) at time *t.sub.1*. The snapshot may be received from a source system, such as source system **152**. The snapshot tree in conjunction with a plurality of object metadata structures may provide a complete view of the source system associated with the snapshot for the particular point in time. In some embodiments, the snapshot tree in conjunction with a plurality of object file metadata structures may provide a complete view of the object associated with the snapshot for the particular point in time.

(120) A root node is the starting point of a snapshot tree and may include pointers to one or more other nodes. An intermediate node is a node to which another node points (e.g., root node, other intermediate node) and includes one or more pointers to one or more other nodes. A leaf node is a node at the bottom of a snapshot tree. Each node of the tree structure includes a view identifier of a view with which the node is associated (e.g., TreeID).

(121) A leaf node may be configured to store key-value pairs. A data key *k* is a lookup value by which a particular leaf node may be accessed. For example, “1” is a data key that may be used to look up “DATA1” of leaf node **322**. The data key *k* may correspond to a brick number of a data brick. A data brick may be comprised of one or more data blocks. In some embodiments, the leaf node is configured to store a brick identifier, file system metadata, object metadata, file size, directory structure, file permissions, physical storage locations of the files, etc. A leaf node may store a data key *k* and a pointer to a location that stores the value associated with the data key. A leaf node may store a pointer to a root node of a metadata structure

(e.g., object metadata structure or an object file metadata structure).

(122) A root node or an intermediate node may include one or more node keys. The node key may be an integer value or a non-integer value. Each node key indicates a division between the branches of the node and indicates how to traverse the tree data structure to find a leaf node, i.e., which pointer to follow. For example, root node **302** may include a node key of “3.” A data key *k* of a key-value pair that is less than or equal to the node key is associated with a first branch of the node and a data key *k* of a key-value pair that is greater than the node key is associated with a second branch of the node. In the above example, to find a leaf node storing a value associated with a data key of “1,” “2,” or “3,” the first branch of root node **302** may be traversed to intermediate node **312** because the data keys of “1,” “2,” and “3” are less than or equal to the node key “3.” To find a leaf node storing a value associated with a data key of “4” or “5,” the second branch of root node **302** may be traversed to intermediate node **314** because data keys “4” and “5” are greater than the node key of “3.”

(123) A data key *k* of a key-value pair is not limited to a numerical value. In some embodiments, non-numerical data keys may be used for a data key-value pair (e.g., “name,” “age,” etc.) and a numerical number may be associated with the non-numerical data key. In some embodiments, a hash function may be associated with the non-numerical data key. The hash function may determine which branch of a node with which the non-numerical data key is associated. The hash function is applied to a data key to provide a hash. The hash may determine which branch of a node with which the non-numerical data key is associated.

(124) In the example shown, root node **302** includes a pointer to intermediate node **312** and a pointer to intermediate node **314**. Root node **302** includes a NodeID of “R1” and a TreeID of “1.” The NodeID identifies the name of the node. The TreeID identifies the view with which the node is associated. When a change is made to data stored in a leaf node as described with respect to FIGS. 3B, 3C, and 3D, the TreeID is used to determine whether a copy of a node is to be made.

(125) Root node **302** includes a node key that divides a set of pointers into two different subsets. Leaf nodes (e.g., “1-3”) with a data key *k* that is less than or equal to the node key are associated with a first branch and leaf nodes (e.g., “4-5”) with a data key *k* that is greater than the node key are associated with a second branch. Leaf nodes with a data key of “1,” “2,” or “3” may be found by traversing tree data structure **300** from root node **302** to intermediate node **312** because the data keys have a value that is less than or equal to the node key. Leaf nodes with a data key of “4” or “5” may be found by traversing tree data structure **300** from root node **302** to intermediate node **314** because the data keys have a value that is greater than the node key.

(126) Root node **302** includes a first set of pointers. The first set of pointers associated with a data key less than the node key (e.g., “1,” “2,” or “3”) indicates that traversing tree data structure **300** from root node **302** to intermediate node **312** may lead to a leaf node with a data key of “1,” “2,” or “3.” Intermediate node **314** includes a second set of pointers. The second set of pointers associated with a data key greater than the node key indicates that traversing tree data structure **300** from root node **302** to intermediate node **314** may lead to a leaf node with a data key of “4” or “5.”

(127) Intermediate node **312** includes corresponding pointers to leaf nodes **322**, **324**, **326**. Intermediate node **314** includes a NodeID of “I1” and a TreeID of “1.” Intermediate node **312** includes a first node key of “1” and a second node key of “2.” The data key *k* for leaf node **322** is a value that is less than or equal to the first node key. The data key *k* for leaf node **324** is a value that is greater than the first node key and less than or equal to the second node key. The data key *k* for leaf node **326** is a value that is greater than the second node key. The pointer to leaf node **322** indicates that traversing tree data structure **300** from intermediate node **312** to leaf node **322** may lead to the node with a data key of “1.” The pointer to leaf node **324** indicates that traversing tree data structure **300** from intermediate node **312** to leaf node **324** may lead to the node with a data key of “2.” The pointer to leaf node **326** indicates that traversing tree data structure **300** from intermediate node **312** to leaf node **326** may lead to the node with a data key of “3.”

(128) Intermediate node **314** includes a pointer to leaf node **328** and a pointer to leaf node **330**. Intermediate node **312** includes a NodeID of “I2” and a TreeID of “1.” Intermediate node **314** includes a node key of “4.” The data key *k* for leaf node **328** is a value that is less than or equal to the node key. The data key *k* for leaf node **330** is a value that is greater than the node key. The pointer to leaf node **328** indicates that traversing tree data structure **300** from intermediate node **314** to leaf node **328** may lead to the node with a data key of “4.” The pointer to leaf node **330** indicates that traversing tree data structure **300** from intermediate node **314** to leaf node **330** may lead to the node with a data key of “5.”

(129) Leaf nodes **322**, **324**, **326**, **328**, **330** include respective data key-value pairs of “1: DATA1,” “2: DATA2,” “3: DATA3,” “4: DATA4,” “5: DATA5.” Leaf nodes **322**, **324**, **326**, **328**, **330** include respective Node IDs of “L1,” “L2,” “L3,” “L4,” “L5.” Leaf nodes **322**, **324**, **326**, **328**, **330** all include a TreeID of “1.” In some embodiments, leaf nodes **322**, **324**, **326**, **328**, or **330** are configured to store metadata. In other embodiments, leaf nodes **322**, **324**, **326**, **328**, or **330** are configured to store content file data when the size of the content file is less than or equal to a limit size. In other embodiments, leaf nodes **322**, **324**, **326**, **328**, or **330** are modes and are configured to store a pointer to or an identifier of a respective object metadata structure (e.g., object metadata tree).

(130) FIG. 3B is a block diagram illustrating an embodiment of a cloned snapshot tree. A CDP service may store a log that includes one or more entries. An entry may include a snapshot marker that indicates a reference restoration point is to be generated. A file system manager may detect that an entry includes a snapshot marker. In response to detecting the snapshot marker, the file system manager may be configured to generate a reference restoration point.

(131) A snapshot tree associated with a previous reference restoration point may be cloned to generate the reference restoration point. A snapshot tree may be cloned when a snapshot tree is added to a tree data structure. In some embodiments, tree data structure **350** may be created by a backup system, such as backup system **162**. A CDP service of the backup system may store a log for a plurality of intercepted write IOs. The log may be comprised of a plurality of entries. One of the entries may include a snapshot marker, which indicates that a reference restoration point is to be generated. A file system manager of the backup system may detect the snapshot marker. Generating a reference restoration point includes cloning a snapshot tree associated with a previous reference restoration point. In some embodiments, the previous reference restoration point is a full reference restoration point. In some embodiments, the previous reference restoration point is an intermediate reference restoration point.

(132) In the example shown, tree data structure **350** includes root nodes **302**, **304**, intermediate nodes **312**, **314**, and leaf nodes **322**, **324**, **326**, **328**, and **330**. Tree data structure **350** may correspond to a version of file system data at a particular moment in time. In some embodiments, tree data structure **350** corresponds to a version of an object at a particular moment in time. The tree data structure can be used to capture different versions of file system data/objects at different moments in time. The tree data structure allows a chain of snapshot trees to be linked together by allowing a node of a later version of a snapshot tree to reference a node of a previous version of a snapshot tree. For example, a snapshot tree with root node **304** is linked to a snapshot tree with root node **302**. Each time a reference restoration point is generated, a new root node may be created and the new root node includes the same set of pointers included in the previous root node, that is, the new root node of the snapshot tree may be linked to one or more intermediate nodes associated with a previous snapshot tree. The new root node also includes a different NodeID and a different TreeID. The TreeID may indicate a snapshot to which a node is associated.

(133) In some embodiments, a root node is associated with a current view. A current view may still accept one or more changes to the data. The TreeID of a root node indicates a snapshot with which the root node is associated. For example, root node **302** with a TreeID of “1” is associated with a first reference restoration point and root node **304** with a TreeID of “2” is associated with a second reference restoration point. In the example shown, root node **304** is associated with a current view.

(134) In other embodiments, a root node is associated with a snapshot view. A snapshot view may represent a state of the file system data/object data at a particular moment in time in the past and is not updated. In the example shown, root node **302** is associated with a snapshot view of the file system data.

(135) In the example shown, root node **304** is a copy of root node **302**. Similar to root node **302**, root node **304** includes the same pointers as root node **302**. However, root node **304** includes a NodeID of “R2” and a TreeID of “2.”

(136) FIG. 3C is a block diagram illustrating an embodiment of modifying a cloned snapshot tree. A cloned snapshot tree may be modified as part of the process to generate a reference restoration point. A CDP service may store a log that includes one or more entries. An entry may be associated with an IO that is intercepted from an object. The IOs associated with the one or more entries may be applied to the cloned snapshot tree to generate a view corresponding to a reference restoration point. In some embodiments, the IOs associated with the one or more entries are applied to the cloned snapshot tree to generate a view corresponding to a specified recovery time point. In some embodiments, the IOs associated with the one or more entries are applied to the cloned snapshot tree to generate a view corresponding to a current state of a

cloned hot standby.

(137) In the example shown, tree data structure **355** may be modified by a file system manager, such as file system manager **167**. A snapshot tree with a root node **304** may be a current view of the file system data or object data. A current view represents a state of the file system data or object that is in the process of being brought up-to-date and capable of receiving one or more modifications to the snapshot tree that correspond to modifications to the file system data or object. Because a snapshot represents a perspective of the file system data or object that is “frozen” in time, one or more copies of one or more nodes affected by a change to file system data are made.

(138) The snapshot tree with root node **304** may correspond to a reference restoration point at time $t_{sub.10}$. A file system manager may access a log stored by a CDP service and apply to the cloned snapshot tree, the IOs that occurred after the last reference restoration point and before the snapshot marker associated with a current reference restoration point.

(139) In the example shown, a log entry indicates that the value “DATA4” has been modified to be “DATA4’.” In some embodiments, the value of a key value pair has been modified. In other embodiments, the value of the key value pair is the data associated with a content file that is smaller than or equal to a limit size. In other embodiments, the value of the key value pair points to a different metadata structure. In some embodiments, the different metadata structure corresponds to a modified version of an object. In some embodiments, the different metadata structure corresponds to a modified version of an object file. The different metadata structure may be a modified version of the metadata structure to which the leaf node previously pointed.

(140) The file system manager starts at root node **304** because that is the root node associated with the cloned snapshot tree. The value “DATA4” is associated with the data key “4.” The file system manager traverses snapshot tree **355** from root node **304** until it reaches a target node, in this example, leaf node **328**. The file system manager compares the TreeID at each intermediate node and leaf node with the TreeID of the root node. In the event the TreeID of a node matches the TreeID of the root node, the file system manager proceeds to the next node. In the event the TreeID of a node does not match the TreeID of the root node, a shadow copy of the node with the non-matching TreeID is made. A shadow copy is a node copy that includes the same pointers as the copied node, but includes a different NodeID and a different TreeID. For example, to reach a leaf node with a data key of “4,” the file system manager begins at root node **304** and proceeds to intermediate node **314**. The file system manager compares the TreeID of intermediate node **314** with the TreeID of root node **304**, determines that the TreeID of intermediate node **314** does not match the TreeID of root node **304**, and creates a copy of intermediate node **314**. The intermediate node **316** includes the same set of pointers as intermediate node **314**, but includes a TreeID of “2” to match the TreeID of root node **204**. Intermediate node **316** is a shadow copy of intermediate node **314**.

(141) The file system manager updates a pointer of root node **304** to point to intermediate node **316** instead of pointing to intermediate node **314**. The file system manager traverses tree data structure **355** from intermediate node **316** to leaf node **328**, determines that the TreeID of leaf node **328** does not match the TreeID of root node **304**, and creates a copy of leaf node **328**. Leaf node copy **332** stores the modified value “DATA4’” and includes the same TreeID as root node **304**. The file system manager updates a pointer of intermediate node **316** to point to leaf node **332** instead of pointing to leaf node **328**.

(142) In some embodiments, leaf node **332** stores the value of a key value pair that has been modified. In other embodiments, leaf node **332** stores the modified data associated with a content file that is smaller than or equal to a limit size. In other embodiments, leaf node **332** stores a pointer to a metadata structure corresponding to the modified object. In other embodiments, leaf node **332** stores a pointer to a metadata tree corresponding to the modified object file.

(143) FIG. 3D is a block diagram illustrating an embodiment of a modified cloned snapshot tree. Tree data structure **355** shown in FIG. 3D illustrates a result of the modifications made to a snapshot tree as described with respect to FIG. 3C. The snapshot tree with root node **302** may correspond to a first reference restoration point and the snapshot tree with root node **304** may correspond to a second reference restoration point. In this example, both snapshot trees are fully hydrated restoration points. A fully hydrated restoration point corresponds to a state which may be restored without having to apply one or more logged IOs. For example, a source system or an object may be restored to a state associated with time $t_{sub.1}$ by traversing the snapshot tree having root node **302**. A source system or an object may be restored to a state associated with time $t_{sub.10}$ by traversing the snapshot tree having root node **304**. A source system or an object may

be restored to a state between time t.sub.1 and t.sub.10 by applying one or more data changes that are stored in a log file to the snapshot tree having root node **302** and then traversing the modified snapshot tree.

(144) FIG. **4A** is a block diagram illustrating an embodiment of a tree data structure. In some embodiments, tree data structure **400** may be created by a backup system, such as backup system **162**. In the example shown, tree data structure **400** corresponds to an object and stores the metadata associated with the object. A leaf node of a snapshot tree associated with file system data, such as a leaf node of tree data structures **300**, **350**, **355**, may include a pointer to a tree data structure corresponding to an object, such as tree data structure **300**. Tree data structure **400** may correspond to a reference backup of an object, e.g., a full backup of an object.

(145) In other embodiments, tree data structure **400** corresponds to an object file and stores the metadata associated with the object file. A leaf node of a snapshot tree associated with an object, such as a leaf node of tree data structures **300**, **350**, **355**, may include a pointer to a tree data structure corresponding to an object file, such as tree data structure **300**.

(146) A tree data structure corresponding to an object or object file is a metadata structure, but is used to organize the data blocks associated with an object or object file that are stored on the backup system. Tree data structure **400** may be referred to as a “metadata structure” or a “metadata tree.”

(147) A tree data structure corresponding to an object or object file at a particular point in time (e.g., a particular version) may be comprised of a root node, one or more levels of one or more intermediate nodes, and one or more leaf nodes. In some embodiments, a tree data structure corresponding to an object or object file is comprised of a root node and one or more leaf nodes without any intermediate nodes. Tree data structure **400** may be a reference restoration point of an object or object file at a particular point in time t, for example at time t.sub.1. A tree data structure associated with file system data may include one or more pointers to one or more tree data structures corresponding to one or more objects. For example, a tree data structure may be associated with the file system data of source system **152** and include one or more pointers to tree data structures corresponding to the one or more objects **153**.

(148) In the example shown, tree data structure **400** includes an object root node **402**, object intermediate nodes **412**, **414**, and object leaf nodes **422**, **424**, **426**, **428**, **430**. Although tree data structure **400** includes one intermediate level between root node **402** and leaf nodes **422**, **424**, **426**, **428**, **430**, any number of intermediate levels may be implemented. Similar to the snapshot trees described above, each node includes a “NodeID” that identifies the node and a “TreeID” that identifies a view with which the node is associated.

(149) In the example shown, root node **402** includes a pointer to intermediate node **412** and a pointer to intermediate node **414**. Root node **402** includes a NodeID of “OR1” and a TreeID of “1.”

(150) In the example shown, intermediate node **412** includes corresponding pointers to leaf nodes **422**, **424**, **426**. Intermediate node **412** includes a NodeID of “011” and a TreeID of “1.” Intermediate node **412** includes a first node key and a second node key. The data key k for leaf node **422** is a value that is less than or equal to the first node key. The data key for leaf node **424** is a value that is greater than the first node key and less than or equal to the second node key. The data key for leaf node **426** is a value that is greater than the second node key. The pointer to leaf node **422** indicates that traversing tree data structure **400** from intermediate node **412** to leaf node **422** may lead to the node with a data key of “1.” The pointer to leaf node **424** indicates that traversing tree data structure **400** from intermediate node **412** to leaf node **424** may lead to the node with a data key of “2.” The pointer to leaf node **426** indicates that traversing tree data structure **400** from intermediate node **412** to leaf node **426** may lead to the node with a data key of “3.”

(151) In the example shown, intermediate node **414** includes a pointer to leaf node **428** and a pointer to leaf node **430**. Intermediate node **414** includes a NodeID of “012” and a TreeID of “1.” Intermediate node **414** includes a node key. The data key k for leaf node **428** is a value that is less than or equal to the node key. The data key for leaf node **430** is a value that is greater than the node key. The pointer to leaf node **428** indicates that traversing tree data structure **400** from intermediate node **414** to leaf node **428** may lead to the node with a data key of “4.” The pointer to leaf node **430** indicates that traversing tree data structure **400** from intermediate node **414** to leaf node **430** may lead to the node with a data key of “5.”

(152) Leaf node **422** includes a data key-value pair of “1: Brick 1.” “Brick 1” is a brick identifier that identifies the data brick that is associated with one or more data chunks of an object or object file. Leaf node **422** includes NodeID of “OL1” and a TreeID of “1.” To view the value associated with a data key of “1,” tree data structure **400** is traversed from root node **402** to intermediate node **412** to leaf node **422**.

(153) Leaf node **424** includes a data key-value pair of “2: Brick 2.” “Brick 2” is a brick identifier that

identifies the data brick that is associated with one or more data chunks of an object or object file. Leaf node **424** includes NodeID of “OL2” and a TreeID of “1.” To view the value associated with a data key of “2,” tree data structure **400** is traversed from root node **402** to intermediate node **412** to leaf node **424**.

(154) Leaf node **426** includes a data key-value pair of “3: Brick 3.” “Brick 3” is a brick identifier that identifies the data brick that is associated with one or more data chunks of an object or object file. Leaf node **426** includes NodeID of “OL3” and a TreeID of “1.” To view the value associated with a data key of “3,” tree data structure **400** is traversed from root node **402** to intermediate node **412** to leaf node **426**.

(155) Leaf node **428** includes a data key-value pair of “4: Brick 4.” “Brick 4” is a brick identifier that identifies the data brick that is associated with one or more data chunks of an object or object file. Leaf node **428** includes NodeID of “OL4” and a TreeID of “1.” To view the value associated with a data key of “4,” tree data structure **400** is traversed from root node **402** to intermediate node **414** to leaf node **428**.

(156) Leaf node **430** includes a data key-value pair of “5: Brick 5.” “Brick 5” is a brick identifier that identifies the data brick that is associated with one or more data chunks of an object or object file. Leaf node **430** includes NodeID of “OL5” and a TreeID of “1.” To view the value associated with a data key of “5,” tree data structure **400** is traversed from root node **402** to intermediate node **414** to leaf node **430**.

(157) An object or object file may be comprised of a plurality of data chunks. A brick may be associated with one or more data chunks. In the example shown, leaf nodes **422**, **424**, **426**, **428**, **430** each store a corresponding brick identifier. The location of the one or more data chunks associated with a data brick may be identified using one or more data structures (e.g., list, table, etc.) stored in metadata store **164**. A first data structure (e.g., chunk metadata table) may store information that associates a brick identifier with one or more chunk identifiers and one or more chunk file identifiers. A second data structure (e.g., chunk file metadata table) may associate a chunk file identifier with a chunk file storing a plurality of data chunks. In some embodiments, the first data structure and the second data structure are combined as a single data structure. The one or more data chunks associated with a data brick may be located based on the chunk metadata table and the chunk file metadata table. For example, a first data brick having a first brick identifier may be associated with a first chunk identifier (e.g., SHA-1 hash value).

(158) The first chunk identifier may be used in conjunction with the chunk metadata table to identify a chunk file identifier. A chunk file having the identified chunk file identifier is comprised of a plurality of data chunks. The chunk file metadata table may be used to identify a location of the plurality of data chunks. The chunk file metadata table may include offset information of the plurality of data chunks within a chunk file.

(159) FIG. **4B** is a block diagram illustrating an embodiment of a cloned metadata structure. A CDP service may store a log that includes one or more entries. An entry may include a snapshot marker that indicates a reference restoration point is to be generated. A file system manager may detect that an entry includes a snapshot marker. In response to detecting the snapshot marker, the file system manager may be configured to generate a reference restoration point.

(160) A snapshot tree corresponding to an object or object file may be cloned to as part of a process to generate a reference restoration point. In some embodiments, tree data structure **450** may be created by a backup system, such as backup system **162**. A CDP service of the backup system may store a log for a plurality of intercepted write IOs. The log may be comprised of a plurality of entries associated with an object or object file. One of the entries may include a snapshot marker, which indicates that a reference restoration point is to be generated. A file system manager of the backup system may detect the snapshot marker. Generating a reference restoration point includes cloning a snapshot tree associated with a previous reference restoration point. In some embodiments, the previous reference restoration point corresponds to a full reference restoration point of an object or object file. In some embodiments, the previous reference restoration point corresponds to an intermediate reference restoration point.

(161) A tree data structure corresponding to an object or object file is a snapshot tree, but stores metadata associated with the object or object file. The tree data structure corresponding to an object or object file can be used to capture different versions of the object or object file at different moments in time. In some embodiments, the tree data structure allows a chain of metadata structures corresponding to different versions of an object or object file to be linked together by allowing a node of a later version of a metadata structure to reference a node of a previous version of a metadata structure. A metadata structure is comprised of a root node, one or more levels of one or more intermediate nodes, and one or more leaf nodes.

(162) A root node or an intermediate node of a version of a metadata structure may reference an intermediate node or a leaf node of a previous version of a metadata structure. Similar to the snapshot tree, the metadata structure allows different versions of object data to share nodes and allows data changes to an object or object file to be tracked.

(163) In the example shown, tree data structure **450** includes a first metadata structure comprising root node **402**, intermediate nodes **412**, **414**, and leaf nodes **422**, **424**, **426**, **428**, and **430**. Tree data structure **450** also includes a second metadata structure that may be a snapshot of an object or object file at a particular point in time $t+n$, for example at time $t_{\text{sub.10}}$. The second metadata structure is comprised of root node **404**, intermediate nodes **412**, **414**, and leaf nodes **422**, **424**, **426**, **428**, and **430**. To create a reference restoration point of an object or object file, a new root node is created. The new root node is a copy of the root node associated with the previous reference restoration point of the object or object file and includes the same set of pointers as the copied root node.

(164) In the example shown, root node **404** includes a set of pointers to intermediate nodes **412**, **414**, which are intermediate nodes associated with a previous reference restoration point. In some embodiments, root node **404** is associated with a current view. The current view may represent a state of the object or object file that is in the process of being brought up-to-date and is capable of receiving one or more modifications to the metadata structure that correspond to modifications to the file data. The TreeID of a root node indicates a reference restoration point with which the root node is associated. For example, root node **402** with a TreeID of “1” is associated with a first reference restoration point and root node **404** with a TreeID of “2” is associated with a second reference restoration point. In other embodiments, root node **404** is associated with a snapshot view. A snapshot view may represent a state of the object or object file at a particular moment in time in the past and is not updated.

(165) In the example shown, root node **404** is a copy of root node **402**. Root node **404** includes a first set of pointers to intermediate node **412**. The first set of pointers associated with a data key (e.g., “1,” “2,” or “3”) less than or equal to the node key indicates that traversing a file metadata tree included in tree data structure **450** from root node **404** to intermediate node **412** may lead to a leaf node with a data key of “1,” “2,” or “3.” Root node **404** includes a second set of pointers to intermediate node **414**. The second set of pointers associated with a data key greater than the node key indicates that traversing a metadata structure included in tree data structure **450** from root node **404** to intermediate node **414** may lead to a leaf node with a data key of “4” or “5.”

(166) FIG. 4C is a block diagram illustrating an embodiment of modifying a cloned metadata structure. A cloned metadata structure may be generated to generate an application consistent recovery point. A CDP service may store a log that includes one or more entries. An entry may be associated with an IO that is intercepted from an object. The IOs associated with the one or more entries may be applied to the cloned metadata structure to generate the reference restoration point of an object or object file. In some embodiments, the IOs associated with the one or more entries are applied to the cloned metadata structure to generate a view corresponding to a specified recovery time point of an object or object file. In some embodiments, the IOs associated with the one or more entries are applied to the cloned metadata structure to generate a view corresponding to a current state of a cloned hot standby that corresponds to an object or object file.

(167) In the example shown, tree data structure **480** may be modified by a file system manager, such as file system manager **167**. A metadata structure with root node **404** may be a current view. A current view may represent a state of an object or object file that is in the process of being brought up-to-date and capable of receiving one or more modifications to the metadata structure tree that correspond to modifications to the object data.

(168) In some embodiments, the object data may be modified such that one of the data chunks is replaced by another data chunk. When a data chunk of object data associated with a previous reference restoration point is replaced with a new data chunk, the data brick associated with the new data chunk may be different. A leaf node of a metadata structure stores a brick identifier associated with a particular brick associated with the data chunk. To represent this modification to the file data, a corresponding modification is made to a current view of a metadata structure. The current view of the metadata structure is modified because the previous metadata structure is a snapshot view and can no longer be modified. The data chunk of the object data that was replaced has a corresponding leaf node in the previous metadata structure. A new leaf node in the current view of the metadata structure is created, as described herein, that corresponds to the new data

chunk. The new leaf node includes an identifier associated with the new reference restoration point. The new leaf node may also store the chunk identifier associated with the modified data chunk.

(169) In the example shown, a data chunk included in “Brick 4” has been modified. The data chunk associated with “Brick 4” has been replaced with a data chunk associated with “Brick 6.” The file system manager starts at root node **404** because that is the root node corresponding to the reference restoration point at time t.sub.10. The value “Brick 4” is associated with the data key “4.” The file system manager traverses tree data structure **480** from root node **404** until it reaches a target node, in this example, leaf node **428**. The file system manager compares the TreeID at each intermediate node and leaf node with the TreeID of the root node. In the event the TreeID of a node matches the TreeID of the root node, the file system manager proceeds to the next node. In the event the TreeID of a node does not match the TreeID of the root node, a shadow copy of the node with the non-matching TreeID is made. For example, to reach a leaf node with a data key of “4,” the file system manager begins at root node **404** and proceeds to intermediate node **414**. The file system manager compares the TreeID of intermediate node **414** with the TreeID of root node **404**, determines that the TreeID of intermediate node **414** does not match the TreeID of root node **404**, and creates a copy of intermediate node **414**. Intermediate node copy **416** is a shadow copy and includes the same set of pointers as intermediate node **414**, but includes a TreeID of “2” to match the TreeID of root node **404**. The file system manager updates a pointer of root node **404** to point to intermediate node **416** instead of pointing to intermediate node **414**. The file system manager traverses tree data structure **480** from intermediate node **416** to leaf node **428**, determines that the TreeID of leaf node **428** does not match the TreeID of root node **404**, and creates a copy of leaf node **428**. Leaf node **432** is a copy of leaf node **428**, but stores the brick identifier “Brick 6” and includes the same TreeID as root node **404**. The file system manager updates a pointer of intermediate node **416** to point to leaf node **432** instead of pointing to leaf node **428**.

(170) FIG. 4D is a block diagram illustrating an embodiment of a modified cloned metadata structure. The metadata tree **480** shown in FIG. 4D illustrates a result of the modifications made to file metadata tree **480** as described with respect to FIG. 4C.

(171) The metadata structure with root node **402** may correspond to a first reference restoration point of an object or object file and the snapshot tree with root node **404** may correspond to a second reference restoration point of an object or object file. In this example, both metadata structures are fully hydrated restoration points. A fully hydrated restoration point corresponds to a state which may be restored without having to apply one or more logged IOs. For example, an object or object file may be restored to a state associated with time t.sub.1 by traversing the metadata structure having root node **402**. An object or object file may be restored to a state associated with time t.sub.10 by traversing the metadata structure having root node **404**. An object or an object file may be restored to a state between time t.sub.1 and t.sub.10 by applying one or more data changes that are stored in a log file to the metadata structure having root node **402** and then traversing the modified snapshot tree.

(172) FIG. 5 is a flow chart of a process for providing continuous data protection in accordance with some embodiments. In the example shown, process **500** may be implemented by a backup system, such as backup system **162**.

(173) At **502**, a reference snapshot of a storage is stored. A reference snapshot may correspond to a full backup snapshot. In some embodiments, a full backup snapshot of a source system is performed. The full backup snapshot may include all of the file system data associated with the source system (e.g., content data, metadata). In some embodiments, a full backup snapshot of an object is performed. Examples of objects include, but are not limited to, a drive, a volume, a VM, a database, an application, a container, a pod, etc. The full backup snapshot may include all of the object data associated with an object (e.g., content data, metadata).

(174) In some embodiments, the reference snapshot of the storage corresponds to an initial full backup snapshot of the source system. In some embodiments, the reference snapshot of the storage corresponds to an initial full backup snapshot of an object. In other embodiments, the reference snapshot of the storage corresponds to a subsequent full backup snapshot of the source system. In other embodiments, the reference snapshot of the storage corresponds to a subsequent full backup snapshot of the object. A subsequent full backup snapshot may be performed in the event a gap in a sequence of IOs is detected.

(175) A backup system may ingest the data included in the reference snapshot, perform a backup of the data included in the reference snapshot, and generate a tree data structure based on the data included in the

reference snapshot. The tree data structure may include a snapshot tree and one or more metadata structures. In some embodiments, the one or more metadata structures correspond to one or more objects hosted on a source system. In some embodiments, the one or more metadata structures correspond to one or more object files associated with the object. The tree data structure corresponding to the reference snapshot may be a full reference restoration point.

(176) At **504**, one or more data changes that modify the storage are received and stored. Each of the CDP-enabled objects hosted on a source system may have an associated write filter. A write filter may be configured to intercept IOs associated with a CDP-enabled object and filter out the read IOs. In some embodiments, the write filter is configured to forward the intercepted write IOs to a filter daemon hosted on the source system.

(177) The filter daemon may be configured to collect the intercepted write IOs from one or more write filters and stream the intercepted write IOs to a CDP service of the backup system. The filter daemon may be configured to stream the intercepted write IOs as they are received. In some embodiments, the filter daemon is configured to stream batches of intercepted write IOs. The filter daemon may be configured to stream intercepted write IOs to CDP services that correspond to a CDP-enabled object from which the intercepted write IO is received. For example, the filter daemon may receive an intercepted write IO from a first CDP-enabled object and an intercepted write IO from a second CDP-enabled object. The filter daemon may forward the intercepted write IO from the first CDP-enabled object to a first CDP service running on a first cluster node of the backup system and forward the intercepted write IO from the second CDP-enabled object to a second CDP service running on the second cluster node. A CDP service hosted on a cluster node may receive a stream of IOs associated with a CDP-enabled object.

(178) At **506**, the received data changes are logged. A CDP service may have a corresponding storage and store a record of the received stream of IOs in a log (e.g., write-ahead log). The CDP service may maintain a corresponding log for each of the CDP-enabled objects. The CDP service may receive the IOs associated with a CDP-enabled object as long as the CDP-enabled object is active. The log is comprised of a plurality of entries associated with a CDP-enabled object. Each entry may correspond to an IO received from the filter daemon. An entry may have an associated timestamp, IO, and an associated sequence number. In some embodiments, an entry includes a snapshot marker. In the event an entry includes a snapshot marker, a reference restoration point may be generated.

(179) At **508**, the one or more data changes occurring after an instance time of the reference snapshot are applied to the reference snapshot to generate a first incremental snapshot corresponding to a first intermediate reference restoration point. A file system manager of the backup system may detect that a log entry includes a snapshot marker. In response to detecting the snapshot marker, the file system manager may be configured to generate an intermediate reference restoration point.

(180) The tree data structure corresponding to the previous snapshot, e.g., a reference snapshot, is identified and a root node of the tree data structure corresponding to the previous snapshot is cloned. For example, the root node of the tree data structure corresponding to the previous snapshot may be cloned in a manner that is described above with respect to FIG. 3B.

(181) The entries of intercepted write IOs stored in a log file may correspond to one or more data changes to a CDP-enabled object. The entries that occurred after a timestamp associated with the reference snapshot and before a timestamp associated with the snapshot marker may be applied to the cloned tree data structure. The entries of the log file that correspond to the CDP-enabled object change may be applied to the cloned tree data structure. For example, the data changes may be applied to the cloned tree data structure in a manner that is described above with respect to FIGS. 3C-3D.

(182) The cloned tree data structure with the one or more data changes applied corresponds to the first intermediate restoration point.

(183) At **510**, the one or more data changes occurring after an instance time of the first incremental snapshot are applied to the first incremental snapshot to generate a second incremental snapshot corresponding to a second intermediate reference restoration point.

(184) A second entry included in the log may include a snapshot marker. In response to detecting the snapshot marker in the second entry, the file system manager may be configured to generate a second intermediate reference restoration point.

(185) The tree data structure corresponding to the previous snapshot, e.g., a reference snapshot, is identified and a root node of the tree data structure corresponding to the previous snapshot is cloned. For example, the

root node of the tree data structure corresponding to the previous snapshot may be cloned in a manner that is described above with respect to FIG. 3B.

(186) The entries of intercepted write IOs stored in a log file may correspond to one or more data changes to a CDP-enabled object. The entries that occurred after a timestamp associated with the previous snapshot marker and before a timestamp associated with the current snapshot marker may be applied to the cloned tree data structure. The entries of the log file that correspond to a CDP-enabled object change may be applied to the cloned tree data structure. For example, the data changes may be applied to the cloned tree data structure in a manner that is described above with respect to FIGS. 3C-3D.

(187) The cloned tree data structure with the one or more data changes applied corresponds to a second intermediate reference restoration point.

(188) FIG. 6 is a flow chart illustrating a process of providing a reference restoration point in accordance with some embodiments. In some embodiments, process 600 is implemented by a write filter, such as write filter 105.

(189) At 602, a notification of a snapshot of an object being created is received. The notification may be received from a CDP-enabled object. The CDP-enabled object may include a backup agent that notifies a write filter.

(190) Reference restoration points may be generated based on a backup policy. In some embodiments, the backup policy indicates that reference restoration points are to be generated periodically (e.g., every 15 minutes, every 30 minutes, every 60 minutes, etc.). In some embodiments, the backup policy indicates that reference restoration points are to be generated in response to a user command. In some embodiments, the backup policy indicates that reference restoration points are to be generated after a threshold number of data changes have been logged. In some embodiments, the backup policy indicates that a reference restoration point is to be generated in the event a gap is detected in the sequenced of logged IOs.

(191) In some embodiments, a backup agent hosted on a source system provides to a write filter that is attached to a CDP-enabled object a command to provide a reference restoration point. The source system may include one or more CDP-enabled objects. Each CDP-enabled object may have a corresponding write filter. The backup agent may be configured to provide the command to each of the write filters.

(192) At 604, a snapshot marker is inserted in the stream of data. The snapshot marker may be inserted in the stream of data in between intercepted write IOs. In some embodiments, the backup agent may provide to a CDP-enabled object a command to quiesce the CDP-enabled object for a short period of time (e.g., 1 second). The snapshot marker may correspond to a point in time at which the CDP-enabled object is quiesced. The snapshot marker may include a timestamp that indicates a reference restoration point is associated with CDP-enabled object to which the write filter is attached.

(193) FIG. 7 is a flow chart illustrating a process for providing a reference restoration point in accordance with some embodiments. In some embodiments, process 700 may be implemented by a backup system, such as backup system 162. In some embodiments, process 700 may be implemented to perform some or all of step 508 of process 500. In some embodiments, process 700 may be implemented to perform some or all of step 510 of process 500.

(194) At 702, a snapshot marker is identified in the stream of data changes. A CDP service of a backup system may receive a stream of write IOs. Each write IO included in the stream may be associated with a timestamp and a sequence number. A snapshot marker may be included in the stream of data changes. A snapshot marker may indicate a point in time that is associated with an application consistent recovery point. An application consistent recovery point may capture on-disk data, data in memory, and all transactions in process.

(195) The CDP service may store a log that includes an entry for each of the stream of data changes. The log may store a timestamp, a sequence number, and/or a snapshot marker. A file system manager of the backup system may analyze the log to determine if any of the entries include a snapshot marker.

(196) At 704, it is determined whether the snapshot_ID associated with a snapshot marker is known to the file system manager. A backup system may initiate a snapshot of a CDP-enabled object. The object snapshot may have an associated ID known to the backup system (and the file system manager). A snapshot of the CDP-enabled object may also be initiated by an application associated with the CDP-enabled object. The object snapshot may have an associated ID, but the associated ID is unknown to the backup system.

(197) In the event the file system manager identifies a snapshot marker, the file system manager determines if the snapshot marker is associated with a snapshot ID known to the file system manager. In the event the

snapshot ID is known to the file system manager, process **700** proceeds to **706**. In the event the snapshot ID is unknown to the file system manager, process **700** proceeds to **710**.

(198) At **706**, a tree data structure associated with a previous reference restoration point is cloned. In the event the file system manager determines that an entry includes a snapshot marker, the file system manager may identify a tree data structure associated with the previous reference restoration point and clone the tree data structure associated with the previous reference restoration point. In some embodiments, the tree data structure associated with the previous reference restoration point corresponds to a full reference restoration point. In some embodiments, the tree data structure associated with the previous recovery point correspond to an intermediate reference restoration point.

(199) The tree data structure associated with the previous reference restoration point may be cloned by copying a root node of the tree data structure associated with the previous reference restoration point. The root node copy may be initially configured to point to the same nodes to which the root node of the tree data structure associated with the previous reference restoration point pointed.

(200) At **708**, one or more data changes included in a log file are applied to the cloned tree data structure. A data change included in a log file may correspond to new data associated with a CDP-enabled object, modified data associated with the CDP-enabled object, or deleted data associated with the CDP-enabled object.

(201) In the event a data change included in the log file corresponds to new data associated with the CDP-enabled object, the cloned tree data structure may be updated to include one or more intermediate nodes and one or more leaf nodes corresponding to the new data associated with the CDP-enabled object.

(202) In the event a data change included in the log file corresponds to modified data associated with the CDP-enabled object, the cloned tree data structure may be updated to include one or more intermediate nodes and one or more leaf nodes corresponding to the modified data associated with the CDP-enabled object, and pointers associated with one or more nodes that correspond to the old data may be updated to point to the new data.

(203) In the event a data change included in the log file corresponds to deleted data associated with the CDP-enabled object, the cloned tree data structure may be updated such that pointers associated with one or more nodes that correspond to the deleted data are deleted.

(204) After the one or more data changes included in the log file are applied to the cloned tree data structure, the modified cloned tree data structure corresponds to a reference restoration point.

(205) At **710**, the log is monitored for a snapshot marker.

(206) FIG. **8** is a flow chart illustrating a process for providing continuous data protection in accordance with some embodiments. Process **800** may be implemented by a backup system, such as backup system **162**.

(207) At **802**, a stream of data changes is received. A write filter may be attached to a CDP-enabled object and configured to intercept IOs associated with the CDP-enabled object.

(208) The write filter may be configured to attach a sequence number to the intercepted write TO. The sequence number indicates a relative order of the particular TO in the stream of IOs. The sequence number may include an incarnation ID and a change ID. The incarnation ID may indicate an incarnation of the write filter and may be incremented each time the write filter is restarted after a crash (e.g., object crash, filter crash). In some embodiments, the incarnation ID is not incremented after a clean restart of a CDP-enabled object. The change ID may be incremented for each intercepted write TO.

(209) At **804**, it is determined that a gap exists in the stream of data changes. The sequence number may be represented as [incarnation ID, change ID]. The sequence number may be [0, 0], [0, 1] . . . [0, n]. In the event the write filter goes offline (e.g., a crash) and comes back online, the incarnation ID may be incremented such that the sequence numbers attached to the intercepted write IOs may be [1, 0], [1, 1] . . . [1, n]. A gap may exist in the event the incarnation ID has been incremented. For example, some changes may have occurred during the time the write filter was online.

(210) In other embodiments, the stream of data changes may be missing an TO. For example, the stream of logged IOs may have a sequence of [0, 0], [0, 1], [0, 2], [0, 3], [0, 5]. In this example, an TO with a sequence number of [0, 4] is missing. A gap may exist in the event a change ID has been skipped in the sequence of IOs.

(211) At **806**, a reference snapshot is requested from a storage. A backup system may provide to a backup agent a command to perform a full snapshot. In some embodiments, in response to receiving the command,

the backup agent causes a source system to perform a full snapshot of its entire system. In some embodiments, in response to receiving the command, the backup agent causes a CDP-enabled object associated with the gap to perform a full snapshot. A snapshot marker may be inserted in the stream of IOs associated with the CDP-enabled object. The snapshot marker may indicate that a full reference restoration point was generated. The snapshot marker may include a snapshot ID that is known to the backup system requesting the reference snapshot.

(212) At **808**, the reference snapshot is received. In some embodiments, the reference snapshot includes all of the file system data associated with a source system. In some embodiments, the reference snapshot includes all of the object data associated with a CDP-enabled object.

(213) At **810**, a reference restoration point corresponding to the reference snapshot is generated. A backup system may ingest the data included in the reference snapshot and generate a tree data structure. Tree data structures, such as the examples depicted in FIGS. 3A and 4A may be generated. The tree data structure may include a snapshot tree and one or more metadata structures. In some embodiments, the one or more metadata structures correspond to one or more objects hosted on a source system. In some embodiments, the one or more metadata structures correspond to one or more object files associated with the object.

(214) In some embodiments, the reference restoration point includes all of the file system data associated with the source system. In some embodiments, the reference restoration point includes all of the object data of an object. Examples of objects include, but are not limited to, a VM, a database, an application, a container, a pod, etc.

(215) FIG. 9 is a flow chart illustrating a process for implementing a disaster recovery plan in accordance with some embodiments. In the example shown, process **900** may be implemented by a data management plan orchestrator, such as data management plan orchestrator **114**.

(216) At **902**, a specification of a data protection intent is received. The data protection intent may include an indication of a data pool and one or more data management services (e.g., backup, restore, disaster recovery, replication, migration, analytics, compliance, security, etc.) that are to be applied to the data pool.

(217) A data pool is a description of one or more objects to be included in the data pool. When the one or more data management services are performed, the one or more data management services are performed with respect to the one or objects specified by the data pool. The description may indirectly and/or directly identify the one or more objects. An object may be associated with one or more tags, labels, features, and/or prefixes. A user may specify the one or more objects to be included in the data pool when the one or more data management services are performed based on a tag, label, feature, and/or prefix associated with an object.

(218) The data protection intent specification may include a disaster recovery plan for the one or more objects to be included in the data pool when the disaster recovery plan is performed. The disaster recovery plan may indicate a primary site associated with each of the one or more objects, a disaster recovery site associated with each of the one or more objects, a specified RPO for the one or more objects, and/or a specified RTO for the one or more objects.

(219) At **904**, it is determined whether CDP is needed to achieve the data protection intent. The data management plan orchestrator may analyze the data protection intent specification to determine protection jobs needed for one or more types of objects to be included in the data pool when the disaster recovery plan is performed. The data protection intent specification analysis may identify one or more types of objects specified to be included in the data pool when the disaster recovery plan is performed. For example, the data management plan orchestrator may identify, based on tags, labels, and/or features associated with a particular object, whether the particular object is a high input/output(IO) object (e.g., an object having a number of IOs/min greater than or equal to an IO threshold), whether the particular object is a low IO object (e.g., an object having a number of IOs/min less than or equal to the IO threshold), performance requirements associated with the particular object, etc.

(220) The data management plan orchestrator may determine protection jobs for the one or more objects to be included in the data pool when the disaster recovery plan is performed based on a specified RPO included in the data protection intent specification and the data protection intent specification analysis. In the event CDP is needed to achieve the data protection intent, process **900** proceeds to **906**. In some embodiments, the data management plan orchestrator determines that CDP needs to be enabled to achieve the specified RPO for an object. For example, an object may be a high IO object and CDP needs to be enabled to achieve the specified RPO for the high IO object.

(221) In the event CDP is not needed to achieve the data protection intent, process **900** proceeds to **910**. In some embodiments, the data management plan orchestrator determines that a scheduled backup (e.g., every hour, every six hours, daily, etc.) is sufficient to achieve the specified RPO for an object. For example, an object may be a low IO object and CDP does not need to be enabled to achieve the specified RPO for the low IO object.

(222) At **906**, CDP is enabled for one or more objects to be included in a data pool when a disaster recovery plan is performed. The data management plan orchestrator may provide to a backup system associated with the source system hosting the one or more objects one or more commands to enable CDP for the one or more objects determined to need CDP to achieve the specified RPO. In response to receiving the one or more commands, the backup system may provide to the source system a corresponding write filter for each of the one or more objects determined to need CDP to achieve the specified RPO and enable on the backup system a CDP service to handle write IOs. In response to receiving a write filter, the source system may attach the corresponding write filter to each of the one or more objects determined to need CDP to achieve the specified RPO. In some embodiments, the source system hosting the one or more objects is capable of providing CDP for the one or more objects because the source system is capable of providing its own write filters. The backup system may provide to the source system a command to enable CDP for some or all of the one or more objects. In response to receiving the command, the source system may attach a corresponding write filter to each of the one or more objects determined to need CDP to achieve the specified RPO.

(223) A write filter may be configured to intercept all IOs of an associated object, filter the read IOs, and provide a stream of the intercepted write IOs to the backup system. The write filter may be configured to attach a sequence number to each intercepted write IO. In some embodiments, the backup system is configured to provide an IO filter daemon to the source system. The IO filter daemon of the source system may be configured to receive the intercepted write IOs from each of the write filters of the source system and to provide the intercepted write IOs from each of the write filters of the source system to the backup system. The backup system may include one or more CDP services that are configured to receive one or more intercepted write IOs from the IO filter daemon of the source system. A CDP service may maintain a log for the one or more intercepted write IOs received from the IO filter daemon of the source system.

(224) At **908**, one or more corresponding hot standbys are caused to be generated for each of the one or more CDP-enabled objects included the data pool when the disaster recovery plan is performed. The data management plan orchestrator may send to a backup system and/or a disaster recovery site system one or more commands that causes the corresponding hot standbys to be generated for each of the one or more CDP-enabled objects to be included in the data pool when the disaster plan is performed.

(225) A frequency at which the one or more hot standbys are updated may be based on a determined sync-replication configuration. The data management plan orchestrator may determine a sync-replication configuration for the one or more hot standbys based on a specified RPO included in a data protection intent. If the specified RPO is less than or equal to a RPO threshold duration (e.g., 5 minutes), the sync-replication configuration may require write IOs logged by the CDP service associated with the backup system to be streamed from the backup system to the hot standby as the write IOs are received by the CDP service associated with the backup system from the source system. If the specified RPO is greater than the RPO threshold duration, the sync-replication configuration may require write IOs logged by the CDP service associated with the backup system to be streamed from the backup system to the hot standby at a frequency that allows the specified RPO to be achieved (e.g., every 30 minutes, hourly, etc.). The backup system may batch the write IOs logged by CDP service associated with the backup system and send a batch of write IOs to the disaster recovery site system hosting the hot standby.

(226) A location of the hot standbys may be based on a determined hot-standby configuration. The data management plan orchestrator may determine a hot-standby configuration based on the specified RPO and the specified RTO included in the data protection intent specification, and the data protection intent specification analysis. In some embodiments, the data management plan orchestrator may determine that the specified RPO and the specified RTO may be achieved by generating at a disaster recovery site system specified in the data protection intent specification a corresponding hot standby for each of the one or more objects to be included in the data pool when the disaster recovery plan is performed. In some embodiments, the data management plan orchestrator may determine that the specified RPO and the specified RTO may be achieved by generating on the backup system a corresponding hot standby for each of the one or more

objects to be included in the data pool when the disaster recovery plan is performed. The data management plan orchestrator may establish this hot-standby configuration even though it does not follow the data protection intent specification because the disaster recovery plan specified by the data protection intent specification may be unable to be satisfied or and/or reduces overall system resources needed to setup and achieve the disaster recovery plan specified by the data protection intent specification.

(227) FIG. **10** is a flow chart illustrating a process for generating a hot standby for a continuous data protection-enabled object in accordance with some embodiments. In the example shown, process **1000** may be implemented by a backup system, such as backup systems **144a**, **144b**, **162**, **180**. In some embodiments, process **1000** is implemented to perform some of step **908** of process **900**.

(228) At **1002**, a reference snapshot of a CDP-enabled object stored. A reference snapshot may correspond to a full backup snapshot and include content associated with the CDP-enabled object. The content associated with the CDP-enabled object may include object data. The content associated with the CDP-enabled object is backed up by the backup system. In some embodiments, the reference snapshot of the storage corresponds to an initial full backup snapshot of a CDP-enabled object. In some embodiments, the reference snapshot of the storage corresponds to a subsequent full backup snapshot of the CDP-enabled object. A subsequent full backup snapshot may be performed in the event a gap in a sequence of IOs associated with the CDP-enabled object is detected.

(229) A backup system may ingest the data included in the reference snapshot and generate a tree data structure for the CDP-enabled object. In some embodiments, the tree data structure includes a snapshot tree for the CDP-enabled object and one or more metadata structures that corresponding to one or more object files associated with the CDP-enabled object. In some embodiments, the tree data structure is a metadata structure that corresponds to the CDP-enabled object. The tree data structure corresponding to the reference snapshot may be a reference restoration point for the CDP-enabled object.

(230) At **1004**, one or more instructions to generate a hot standby corresponding to a CDP-enabled object is provided to a determined location for a hot standby. The one or more instructions may include a command to generate a hot standby. The instructions may be provided using one or more APIs associated with an application (e.g., VMWare) that is configured to generate one or more hot standbys. In response to receiving the one or more instructions, a system at the determined location may generate a hot standby placeholder object for the CDP-enabled object. In some embodiments, a hot standby placeholder object is generated for a database. In some embodiments, a hot standby placeholder object is generated for a VM (e.g., VM configuration file that describes a VM).

(231) In some embodiments, the determined location for the hot standby is a disaster recovery site system, such as a source system located at the disaster recovery site or a backup system located at the disaster recovery site. In some embodiments, the determined location for the hot standby corresponding CDP-enabled object is a backup system that is receiving write IOs from a source system hosting the CDP-enabled object.

(232) At **1006**, data included in the reference snapshot is provided to a determined location for a hot standby to generate a hot standby corresponding to the CDP-enabled object. The provided data includes the content associated with the CDP-enabled object. In response to receiving the data, the disaster recovery site system is configured to generate the hot standby corresponding to the CDP-enabled object by populating the hot standby placeholder object based on the received data. The hot standby corresponding to the CDP-enabled object may be used for failover purposes in the event the CDP-enabled object hosted on the source system goes offline.

(233) At **1008**, one or more data changes that modify the CDP-enabled object are received. The one or more data changes include updated content to the CDP-enabled object. The updated content is backed up by the backup system.

(234) A CDP-enabled object may have an associated write filter. A write filter may be configured to intercept IOs associated with the CDP-enabled object and filter out the read IOs. In some embodiments, the write filter is configured to forward the intercepted write IOs to a filter daemon hosted on a source system associated with the CDP-enabled object. The source system may host one or more CDP-enabled objects.

(235) The filter daemon may be configured to collect the intercepted write IOs from one or more write filters and stream the intercepted write IOs to a CDP service of the backup system. A CDP service hosted on a cluster node of the backup system may receive a stream of IOs associated with the object.

(236) At **1010**, the received data changes are logged. The CDP service of the backup system may maintain a

log the one or more intercepted write IOs received from the IO filter daemon of the source system. In some embodiments, the CDP service maintains a corresponding log for each CDP-enabled object hosted by the source system. In some embodiments, the CDP service maintains a log for all of the CDP-enabled objects hosted by the source system.

(237) Each entry of the log may store the metadata associated with an intercepted write IO. For example, the metadata associated with an intercepted write IO may include at least one of the following: sequence number, timestamp, IO offset, IO length, IO data bytes, VM universally unique identifier (UUID), VM disk (VMDK) full path on a host at which the VM is running, etc.

(238) At **1012**, the one or more received data changes are provided to the determined location to update the hot standby corresponding to the CDP-enabled object. A frequency at which the one or more data changes are provided to the determined location may be based a specified RPO included in a data protection intent. If the specified RPO is less than or equal to a RPO threshold duration (e.g., 5 minutes), the one or more received data changes may be streamed from the backup system to the hot standby as the one or more data changes are received by the CDP service associated with the backup system from the source system. If the specified RPO is greater than the RPO threshold duration, the one or more received data changes may be streamed from the backup system to the hot standby at a frequency that allows the specified RPO to be achieved (e.g., every 30 minutes, hourly, etc.). The backup system may batch the one or more data changes and send a batch of one or more data changes to the determined location for the hot standby.

(239) FIG. **11** is a flow chart illustrating a process for generating and maintaining a hot standby in accordance with some embodiments. In the example shown, process **1100** may be implemented by a recovery system, such as source systems **143a**, **143b** or backup systems **144a**, **144b**. In some embodiments, the recovery system is located at a disaster recovery site. In some embodiments, the recovery system is located at a primary site.

(240) At **1102**, one or more instructions to generate a hot standby for a CDP-enabled object is received. In response to receiving the one or more instructions, a hot standby placeholder object may be generated for the CDP-enabled object. In some embodiments, a hot standby placeholder object is generated for a database. In some embodiments, a hot standby placeholder object is generated for a VM.

(241) At **1104**, a hot standby corresponding to the CDP-enabled object is generated based on the received data. The received data may be used to populate the hot standby placeholder object. In some embodiments, a VM hot standby is generated for the CDP-enabled object. The hot standby is a standby version of at least a portion of the CDP-enabled object hosted on the source system. The hot standby is a standby version of at least a portion of the CDP-enabled object hosted on the source system because there is a delay between a point in time when a write IO is applied to a CDP-enabled object hosted on the source system and a point in time when the write IO is applied to the hot standby.

(242) At **1106**, data that corresponds to a CDP-enabled object is received. The data that corresponds to a CDP-enabled object hosted on a source system may be received at a recovery system from a backup system. The data may correspond to a full snapshot of the CDP-enabled object.

(243) At **1108**, one or more data changes for the CDP-enabled object are received. A backup system may receive from a source system hosting a CDP-enabled object the one or more data changes for the CDP-enabled object. The recovery system may receive the one or more data changes from the backup system.

(244) At **1110**, the hot standby corresponding to the CDP-enabled object is updated based on the one or more received data changes. The recovery system may include a CDP service that is configured to apply the one or more received data changes to the hot standby.

(245) FIG. **12** is a flow chart illustrating a process for updating a hot standby in accordance with some embodiments. In the example shown, process **1200** may be implemented by a backup system, such as backup systems **144a**, **144b**, **162**, **180**.

(246) At **1202**, a request to place a hot standby to a specified recovery time point is received. The request may be received from a user via a user device associated with the user. In some embodiments, the request to place (e.g., update) is a request to roll back the hot standby to the specified recovery time point. For example, a current state of a hot standby may be associated with 1:15 μ m and the specified recovery time point may be 12:40 μ m. The hot standby is a standby version of at least a portion of the CDP-enabled object hosted on the source system.

(247) A CDP-enabled object hosted on a source system may need to be rolled back to the specified recovery time point for one or more reasons. For example, the CDP-enabled object may have become infected with

malware, subject to ransomware, critical data may have been permanently deleted, etc. In some embodiments, the request to place is a request to roll forward the hot standby to the specified recovery time point.

(248) At **1204**, the hot standby is cloned. The backup system may send to a recovery system hosting the hot standby one or more instructions to clone the hot standby. In response to receiving the one or more instructions, the recovery system may clone the hot standby. For the example, the hot standby clone may be a VM clone. The backup system may still be streaming write IOs to the hot standby. Cloning the hot standby may pause a state of the hot standby and prevent the streamed write IOs from being applied to the hot standby. For example, write IOs having a timestamp after 1:15 pm may not be applied to the cloned hot standby.

(249) At **1206**, it is determined whether a log entry corresponding to the specified recovery time point exists. Each entry included in a log is associated with a timestamp and a sequence number. The backup system may determine whether the log stores an entry that corresponds to the specified recovery time point. For example, the backup system may determine whether the log stores an entry for 12:40 pm. The entry for 12:40 pm is associated with a sequence number.

(250) In some embodiments, the log entry corresponding to the specified recovery time point is an entry having a timestamp equal to the specified recover time point. In some embodiments, the log entry corresponding to the specified recovery time point is an entry having a timestamp that is before and closest to the specified recovery time point. In some embodiments, the log entry corresponding to the specified recovery time point is an entry having a timestamp that is after and closest to the specified recovery time point.

(251) In the event a log entry corresponding to the specified recovery time point exists, process **1200** proceeds to **1208**. In the event a log entry corresponding to the specified recovery time point does not exist, process **1200** proceeds to **1214**. For example, the specified recovery time point may correspond to a point in time for which a gap in log entries exists.

(252) At **1208**, a view corresponding to a specified recovery time point and a view corresponding to a current state of the cloned hot standby are generated.

(253) The backup system may generate a view that corresponds to the specified recovery time point by identifying a closest reference restoration point that occurs before the particular moment in time. The closest reference restoration point is associated with a corresponding tree data structure. For example, for a specified recovery time point of 12:40 pm, as seen in FIG. 2, the closest reference restoration point is 12:30 pm. The view that corresponds to the specified recovery time point may be generated by cloning the tree data structure that corresponds to the closest reference restoration point, as seen in FIG. 3B or FIG. 4B, and applying to the cloned tree data structure, as seen in FIG. 3C or FIG. 4C, the logged write IOs that occurred from the closest reference restoration point to the specified recovery time point. For example, the logged write IOs that occurred from 12:30 pm-12:40 pm may be applied in sequence (e.g., consecutively) to the cloned tree data structure to generate the view that corresponds to the specified recovery time point. The backup system may stop applying write IOs after it applies the write IO that corresponds to the sequence number associated with 12:40 pm. The view corresponding to the specified recovery time point is a fully hydrated recovery point because it provides a complete view of the CDP-enabled object at the specified recovery time point.

(254) The backup system may generate a view that corresponds to a current state of the cloned hot standby by identifying a closest reference restoration point that occurs before the current state of cloned hot standby. For example, for a current state of a hot standby being associated with 1:15 pm, as seen in FIG. 2, the closest reference restoration point is 1:00 pm. The view that corresponds to the specified recovery time point may be generated by cloning the tree data structure that corresponds to the closest reference restoration point, as seen in FIG. 3B or FIG. 4B, and applying to the cloned tree data structure, as seen in FIG. 3C or FIG. 4C, the logged write IOs that occurred from the closest reference restoration point to the specified recovery time point. For example, the entry for 1:15 pm is associated with a sequence number. The logged write IOs that occurred from 1:00 pm-1:15 pm may be applied to the cloned tree data structure to generate the view that corresponds to the current state of the cloned hot standby. The backup system may stop applying write IOs after it applies the write IO that corresponds to the sequence number associated with 1:15 pm. The view corresponding to a current state of the cloned hot standby is a fully hydrated recovery point because it provides a complete view of the CDP-enabled object at a time corresponding to

the current state.

(255) At **1210**, one or more differences between the view corresponding to the specified recovery time point and the view corresponding to a current state of the cloned hot standby are determined. The one or more differences may be determined by traversing the views to determine one or more non-shared data chunks. One or more of the non-shared data chunks may be associated with the view corresponding to the specified recovery time point. One or more of the non-shared data chunks may be associated with the view corresponding to a current state of the cloned hot standby.

(256) For example, the metadata structure with root node **402** of FIG. **4D** may be the view corresponding to the specified recovery time point and the metadata structure with root node **404** may be the view corresponding to the current state of the cloned hot standby. The one or more non-shared data chunks are the one or more data chunks associated with “Brick 4” and the one or more data chunks associated with “Brick 6.”

(257) At **1212**, one or more data chunks associated with the view corresponding to the specified recovery time point are sent to the recovery system hosting the hot standby. The backup system may send to the recovery system the one or more of the non-shared data chunks that are associated with the view corresponding to the specified recovery time point. For example, the backup system may send the one or more data chunks associated with “Brick 4.” The one or more data chunks are backup data associated with a CDP-enabled object that were previously received from a source system hosting the CDP-enabled object.

(258) In response to receiving the one or more data chunks, the recovery system may apply the one or more received data chunks to the hot standby and update the hot standby to the specified recovery time point.

(259) At **1214**, a notification that the hot standby cannot be rolled back or forward to the specified recovery time point is provided. For example, the hot standby may not be able to roll back or forward to the specified recovery time point because there may be a gap of entries in the log stored by backup system.

(260) FIG. **13** is a flow chart illustrating a process for updating a hot standby in accordance with some embodiments. In the example shown, process **1300** may be implemented by a backup system, such as backup systems **144a**, **144b**, **162**, **180**. Process **1300** may be implemented after a hot standby corresponding to a CDP-enabled object has been generated. In this example, a current state of the hot standby may correspond to a state of the CDP-enabled object at 12:50 pm.

(261) At **1302**, a stream of data changes is received. A write filter may be attached to a CDP-enabled object and configured to intercept IOs associated with the CDP-enabled object. The write filter may be configured to attach a sequence number to the intercepted write IO. The sequence number indicates a relative order of the particular IO in the stream of IOs.

(262) At **1304**, it is determined that a gap exists in the stream of data changes. This may cause the hot standby corresponding to the CDP-enabled object to be in an inconsistent state with respect to the CDP-enabled object hosted on the source system because the one or more missing write IOs may not be applied to the hot standby corresponding to the CDP-enabled object. The hot standby corresponding to the CDP-enabled object may remain in an inconsistent state with respect to the CDP-enabled object hosted on the source system unless gap of write IOs is corrected.

(263) The sequence number may be represented as [incarnation ID, change ID]. The sequence number may be [0, 0], [0, 1] . . . [0, n]. In the event the write filter goes offline (e.g., a crash) and comes back online, the incarnation ID may be incremented such that the sequence numbers attached to the intercepted write IOs may be [1, 0], [1, 1] . . . [1, n]. A gap may exist in the event the incarnation ID has been incremented.

(264) The change ID may be incremented (e.g., 1, 2, 3, . . . , n) for each intercepted write TO. In other embodiments, the stream of data changes may be missing an TO. For example, the stream of logged IOs may have a sequence of [0, 0], [0, 1], [0, 2], [0, 3], [0, 5]. In this example, an TO with a sequence number of [0, 4] is missing. A gap may exist in the event a change ID has been skipped in the sequence of IOs.

(265) In this example, the backup system may detect a gap exists between 12:51 pm and 12:56 pm.

(266) At **1306**, a reference snapshot is requested from a storage. A backup system may provide to a backup agent hosted on a source system a command to perform a full snapshot. In some embodiments, in response to receiving the command, the backup agent causes the source system to perform a full snapshot of its entire system. In some embodiments, in response to receiving the command, the backup agent causes a CDP-enabled object associated with the gap to perform a full snapshot. A snapshot marker may be inserted in the stream of IOs associated with the CDP-enabled object. The snapshot marker may indicate that a full reference restoration point was generated. The snapshot marker may include a snapshot ID that is known to

the backup system requesting the reference snapshot.

(267) In this example, the backup system may request a full snapshot that corresponds a state of the CDP-enabled object at 1:00 pm.

(268) At **1308**, the reference snapshot is received. In some embodiments, the reference snapshot includes all of the file system data associated with a source system. In some embodiments, the reference snapshot includes all of the object data associated with a CDP-enabled object.

(269) At **1310**, the hot standby is cloned. The backup system may send to a recovery system hosting the hot standby one or more instructions to clone the hot standby. In response to receiving the one or more instructions, the recovery system may clone the hot standby. For the example, the hot standby clone may be a VM clone. The backup system may still be streaming write IOs to the hot standby. Cloning the hot standby may pause a state of the hot standby and prevent the streamed write IOs from being applied to the hot standby. For example, a current state of the hot standby may correspond to a state of the CDP-enabled at 12:50 pm.

(270) At **1312**, a view corresponding to the reference restoration point is generated. For example, the backup system may generate a view that corresponds to 1:00 pm. A backup system may generate the view corresponding to the reference restoration point by ingesting the data included in the reference snapshot and generating a tree data structure that corresponds to the reference snapshot. Tree data structures, such as the examples depicted in FIGS. 3A and 4A may be generated. The tree data structure may include a snapshot tree and one or more metadata structures. In some embodiments, the one or more metadata structures correspond to one or more objects hosted on a source system. In some embodiments, the tree data structure includes a snapshot tree that corresponds to the object and one or more metadata structures correspond to one or more object files associated with the object. The view corresponding to the reference restoration point is a fully hydrated recovery point because it provides a complete view of the CDP-enabled object at a time corresponding to the reference restoration point.

(271) At **1314**, a view corresponding to a current state of a cloned hot standby is generated. The backup system may generate a view that corresponds to the current state of the cloned hot standby by identifying a closest reference restoration point that occurs before the current state of cloned hot standby. For example, for a current state of a hot standby being associated with 12:50 pm, as seen in FIG. 2, the closest reference restoration point is 12:45 pm. The view that corresponds to the specified recovery time point may be generated by cloning the tree data structure that corresponds to the closest reference restoration point, as seen in FIG. 3B or FIG. 4B, and applying to the cloned tree data structure, as seen in FIG. 3C or FIG. 4C, the logged write IOs that occurred from the closest reference restoration point to the specified recovery time point. For example, the logged write IOs that occurred from 12:45 pm-12:50 pm may be applied to the cloned tree data structure to generate the view that corresponds to the current state of the cloned hot standby. The view corresponding to a current state of the cloned hot standby is a fully hydrated recovery point because it provides a complete view of the CDP-enabled object at a time corresponding to the current state.

(272) At **1316**, one or more differences between the view corresponding to the reference restoration point and the view corresponding to the current state of the cloned hot standby are determined. The one or more differences may be determined by traversing the views to determine one or more non-shared data chunks. One or more of the non-shared data chunks may be associated with the view corresponding to the current state of the cloned hot standby. One or more of the non-shared data chunks may be associated with the view corresponding to the reference restoration point.

(273) For example, the metadata structure with root node **402** of FIG. 4D may be the view corresponding to the current state of the cloned hot standby and the metadata structure with root node **404** may be the view corresponding to the reference restoration point. The one or more non-shared data chunks that are associated with the view corresponding to the current state of the cloned hot standby may be the one or more data chunks associated with “Brick 4” and the one or more non-shared data chunks that are associated with the view corresponding to the reference restoration point may be the one or more data chunks associated with “Brick 6.”

(274) At **1318**, one or more data chunks associated with the view corresponding to the reference restoration point are sent to the recovery system hosting the hot standby. The backup system may send the one or more of the non-shared data chunks are associated with the view corresponding to the reference restoration point to the recovery system. For example, the backup system may send the one or more data chunks associated

with “Brick 6.” One or more data chunks are backup data associated with a CDP-enabled object that were previously received from a source system hosting the CDP-enabled object.

(275) In response to receiving the one or more data chunks, the recovery system may update the hot standby by applying the one or more received data chunks. As a result, the hot standby corresponding to the CDP-enabled object may be in consistent state with respect to the CDP-enabled object hosted on the source system or in a state that allows the hot standby corresponding to the CDP-enabled object to be updated to be in a consistent state with respect to the CDP-enabled object hosted on the source system.

(276) FIG. 14 is a flow chart illustrating a process for updating a hot standby in accordance with some embodiments. In the example shown, process 1400 may be implemented by a backup system, such as backup systems 144a, 144b, 162, 180. Process 1400 may be implemented after a hot standby corresponding to a CDP-enabled object has been generated.

(277) At 1402, an indication that one or more write IOs were not applied to a hot standby is received. A backup system may receive from a recovery system hosting a hot standby an indication that one or more write IOs included in a stream were not applied to the hot standby because a gap exists in the stream of write IOs. Each IO is associated with a corresponding sequence number. The recovery system may include a CDP service that determines the sequences numbers associated with the stream of write IOs and determines if a gap exists. The recovery system may wait a threshold period of time before determining if the gap exists. Write IOs may be received out-of-sequence but a gap may not exist after waiting the threshold period of time. The recovery system may provide the indication to the backup system in response to determining that a gap exists.

(278) At 1404, a sequence number associated with a last write IO that was applied to the hot standby before the gap is determined. The recovery system hosting the hot standby may determine the last write IO that was applied to the hot standby before the one or more missing write IOs based on a sequence number associated with the last write IO and provide the determined last write IO in the indication that was received by the backup system.

(279) At 1406, a log is replayed from the last write IO before the gap to a most recent IO. The backup system may replay the stream of write IOs from the sequence number associated with the last write IO that was applied to the hot standby and send to the recovery system hosting the hot standby the one or more write IOs that occurred after the last write IO that was applied to the hot standby. In response to receiving the one or more write IOs, the CDP service of the recovery system hosting the hot standby may apply the one or more write IOs to the hot standby.

(280) The invention can be implemented in numerous ways, including as a process; an apparatus; a system; a composition of matter; a computer program product embodied on a computer readable storage medium; and/or a processor, such as a processor configured to execute instructions stored on and/or provided by a memory coupled to the processor. In this specification, these implementations, or any other form that the invention may take, may be referred to as techniques. In general, the order of the steps of disclosed processes may be altered within the scope of the invention. Unless stated otherwise, a component such as a processor or a memory described as being configured to perform a task may be implemented as a general component that is temporarily configured to perform the task at a given time or a specific component that is manufactured to perform the task. As used herein, the term ‘processor’ refers to one or more devices, circuits, and/or processing cores configured to process data, such as computer program instructions.

(281) A detailed description of one or more embodiments of the invention is provided along with accompanying figures that illustrate the principles of the invention. The invention is described in connection with such embodiments, but the invention is not limited to any embodiment. The scope of the invention is limited only by the claims and the invention encompasses numerous alternatives, modifications and equivalents. Numerous specific details are set forth in the description in order to provide a thorough understanding of the invention. These details are provided for the purpose of example and the invention may be practiced according to the claims without some or all of these specific details. For the purpose of clarity, technical material that is known in the technical fields related to the invention has not been described in detail so that the invention is not unnecessarily obscured.

(282) Although the foregoing embodiments have been described in some detail for purposes of clarity of understanding, the invention is not limited to the details provided. There are many alternative ways of implementing the invention. The disclosed embodiments are illustrative and not restrictive.

Claims

1. A method, comprising: receiving a stream of data changes to content of an object associated with a source system; updating a hot standby version of at least a portion of the object maintained at a recovery system by streaming the received stream of data changes to the recovery system, wherein the hot standby version comprises a replica of the object associated with the source system that is maintained to be up to date and ready to use within a specified recovery time objective (RTO); determining that a gap exists in the stream of data changes such that the hot standby version of the at least the portion of the object maintained at the recovery system is in an inconsistent state with respect to the at least the portion of the object associated with the source system, wherein the at least the portion of the object includes continuous data protection (CDP) enabled object; in response to determining that the gap exists in the stream of data changes, obtaining a reference snapshot from the source system; and causing the inconsistent state of the hot standby version of at least the portion of the object maintained at the recovery system to be updated to a state of the portion of the object associated with the reference snapshot.
2. The method of claim 1, wherein a data change included in the stream of data changes is associated with a sequence number.
3. The method of claim 2, wherein the sequence number is associated with an incarnation identifier.
4. The method of claim 3, wherein determining that the gap exists is based on determining that the incarnation identifier has been incremented.
5. The method of claim 2, wherein the sequence number is associated with a change identifier.
6. The method of claim 5, wherein determining that the gap exists is based on determining that the corresponding sequence number associated with the one or more data changes included in the stream of data changes is missing.
7. The method of claim 1, wherein causing the inconsistent state of the hot standby version of at least the portion of the object maintained at the recovery system to be updated to the state of the portion of the object associated with the reference snapshot includes sending to the recovery system one or more instructions to clone the current state of the hot standby version of at least the portion of the object maintained at the recovery system.
8. The method of claim 7, wherein the recovery system clones the inconsistent state of the hot standby version of at least the portion of the object in response to receiving the one or more instructions.
9. The method of claim 1, wherein causing the inconsistent state of the hot standby version of at least the portion of the object maintained at the recovery system to be updated to the state of the portion of the object associated with the reference snapshot includes generating a view corresponding to the object associated with the reference snapshot.
10. The method of claim 1, wherein causing the inconsistent state of the hot standby version of at least the portion of the object maintained at the recovery system to be updated to the state of the portion of the object associated with the reference snapshot includes generating a view corresponding to the current state of the hot standby version of at least the portion of the object.
11. The method of claim 1, wherein causing the inconsistent state of the hot standby version of at least the portion of the object maintained at the recovery system to be updated to the state of the portion of the object associated with the reference snapshot further includes determining one or more differences between a view corresponding to the object associated with the reference snapshot and a view corresponding to the current state of the hot standby version of at least the portion of the object.
12. The method of claim 11, wherein the one or more determined differences correspond to one or more non-shared data chunks.
13. The method of claim 12, wherein causing the inconsistent state of the hot standby version of at least the portion of the object maintained at the recovery system to be updated to the state of the portion of the object associated with the reference snapshot further includes providing the one or more non-shared data chunks to the recovery system.
14. The method of claim 13, wherein in response to receiving the one or more non-shared data chunks, the recovery system updates the hot standby version of at least the portion of the object.
15. A computer program product embodied in a non-transitory computer readable medium and comprising storing computer instructions that, when executed, cause one or more processors to: receiving receive a

stream of data changes to content of an object associated with a source system; update a hot standby version of at least a portion of the object maintained at a recovery system by streaming the received stream of data changes to the recovery system, wherein a hot standby comprises a replica of the object associated with the source system that is up to date and ready to use within a specified recovery time objective (RTO); determine that a gap exists in the stream of data changes such that the hot standby version of the at least the portion of the object maintained at the recovery system is in an inconsistent state with respect to the at least the portion of the object associated with the source system, wherein the at least the portion of the object includes continuous data protection (CDP) enabled object; in response to determining that the gap exists in the stream of data changes, request a reference snapshot from the source system; and causing the inconsistent state of the hot standby version of at least the portion of the object maintained at the recovery system to be updated to a state of the portion of the object associated with the reference snapshot.

16. The computer program product of claim 15, wherein a data change included in the stream of data changes is associated with a sequence number.

17. The computer program product of claim 16, wherein the sequence number is associated with one of or both an incarnation identifier and a change identifier.

18. The computer program product of claim 15, wherein the computer instructions include computer instructions that, when executed, cause the one or more processors to determine that the gap exists based on corresponding sequence numbers associated with the stream of data changes.

19. A system, comprising: one or more processors configured to: receive a stream of data changes to content of an object associated with a source system; update a hot standby version of at least a portion of the object maintained at a recovery system by streaming the received stream of data changes to the recovery system, wherein the hot standby comprises a replica of the object associated with the source system that is up to date and ready to use within a specified recovery time objective (RTO); determine that a gap exists in the stream of data changes such that the hot standby version of the at least the portion of the object maintained at the recovery system is in an inconsistent state with respect to the at least the portion of the object associated with the source system, wherein the at least the portion of the object includes continuous data protection (CDP) enabled object; in response to determining that the gap exists in the stream of data changes, request a reference snapshot from the source system; and causing the inconsistent state of the hot standby version of at least the portion of the object maintained at the recovery system to be updated to a state of the portion of the object associated with the reference snapshot; and a memory coupled to the one or more processors and configured to provide the one or more processors with instructions.
