

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12405861
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Mandal; Vibhavari et al.

Adaptive scaling of topics in a stateful event-driven platform

Abstract

A system facilitating scaling of topics is provided. During operation, the system can suspend a stateful application of a topic and a stateful backup topic on an event-driven platform. The topic corresponds to a stream of events associated with the application. The system can modify configuration at the platform to allow a new set of partitions of the topic and a new set of backup partitions of the backup topic. The system can identify a state of event records from the application in a previous backup partition and identify a new backup partition allocated to the application. If they are different, the system can transfer the state to the new backup partition. The system can resume the application for generating an updated state, which is maintained at the new backup partition, based on the state and subsequent event records associated with the application.

Inventors: Mandal; Vibhavari (Bangalore Karnataka, IN), Sahoo; Sanjaya Kumar (Karnatak, IN), Roy; Tathagata (Karnatak, IN), Alam; Asif (Karnatak, IN), Yadav; Brijesh (Karnatak, IN), Sahoo; Siddhartha (Karnatak, IN)

Applicant: Hewlett Packard Enterprise Development LP (Spring, TX)

Family ID: 92461980

Assignee: Hewlett Packard Enterprise Development LP (Spring, TX)

Appl. No.: 18/655422

Filed: May 06, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240289232 A1	Aug. 29, 2024

Publication Classification

Int. Cl.: G06F12/00 (20060101); G06F11/14 (20060101)

U.S. Cl.:

CPC **G06F11/1461** (20130101); G06F2201/86 (20130101)

Field of Classification Search

CPC: G06F (11/1461); G06F (11/1446); G06F (2201/86); G06F (3/061); G06F (3/0644); G06F (3/0647); G06F (3/065)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6532498	12/2002	Hager	719/318	G06F 9/542
2007/0282979	12/2006	Tuel	709/219	H04L 67/1027

Primary Examiner: Gu; Shawn X

Attorney, Agent or Firm: Hewlett Packard Enterprise Patent Department

Background/Summary

BACKGROUND

(1) Event streaming is the practice of capturing data in real-time from sources in the form of streams of events; storing these event streams for later retrieval; processing and responding to the event streams; and routing the event streams to different destination technologies as needed. Event streaming ensures a continuous flow and interpretation of data so that the right information is at the right place, at the right time.

Description

BRIEF DESCRIPTION OF THE FIGURES

- (1) FIG. 1A illustrates an example of dynamic upscaling of a topic in an event-driven platform, in accordance with an aspect of the present application.
- (2) FIG. 1B illustrates an example of dynamic downscaling of a topic in an event-driven platform, in accordance with an aspect of the present application.
- (3) FIG. 2A illustrates an example of a partition management system increasing the number of partitions of a topic in an event-driven platform, in accordance with an aspect of the present application.
- (4) FIG. 2B illustrates an example of redistributing backup states among the increased number of partitions of a topic in an event-driven platform, in accordance with an aspect of the present application.
- (5) FIG. 3A illustrates an example of a partition management system generating temporary topics for decreasing the number of partitions of a topic in an event-driven platform, in accordance with an aspect of the present application.
- (6) FIG. 3B illustrates an example of a partition management system transferring record logs and states to temporary topics in an event-driven platform, in accordance with an aspect of the present application.

application.

(7) FIG. 3C illustrates an example of a partition management system regenerating topics in an event-driven platform, in accordance with an aspect of the present application.

(8) FIG. 3D illustrates an example of a partition management system transferring record logs and states to regenerated topics in an event-driven platform, in accordance with an aspect of the present application.

(9) FIG. 3E illustrates an example of a partition management system removing temporary topics in an event-driven platform, in accordance with an aspect of the present application.

(10) FIG. 4A presents a flowchart illustrating the process of a partition management system upscaling a topic in an event-driven platform, in accordance with an aspect of the present application.

(11) FIG. 4B presents a flowchart illustrating the process of a stateful application updating a state based on a new backup partition in an event-driven platform, in accordance with an aspect of the present application.

(12) FIG. 4C presents a flowchart illustrating the process of a producer application identifying a new backup partition in an event-driven platform, in accordance with an aspect of the present application.

(13) FIG. 5A presents a flowchart illustrating the process of a partition management system downscaling a topic in an event-driven platform, in accordance with an aspect of the present application.

(14) FIG. 5B presents a flowchart illustrating the process of a stateful application distributing record logs and states to regenerated topics in an event-driven platform, in accordance with an aspect of the present application.

(15) FIG. 5C presents a flowchart illustrating the process of a partition management system resuming consumer applications of a scaled topic in an event-driven platform, in accordance with an aspect of the present application.

(16) FIG. 6 illustrates an example of a computing system facilitating dynamic topic repartitioning in an event-driven platform, in accordance with an aspect of the present application.

(17) FIG. 7 illustrates an example of a computer-readable medium (CRM) that facilitates dynamic topic repartitioning in an event-driven platform, in accordance with an aspect of the present application.

(18) In the figures, like reference numerals refer to the same figure elements.

DETAILED DESCRIPTION

(19) An event-driven platform, such as Apache Kafka, can dynamically incorporate records of events. The event-driven platform can also be referred to as an event-driven platform. An event can be an action detected and recorded by an application or software, such as a click on a website or a customer purchasing an item. The event-driven platform can be distributed on a cluster of servers. A respective type of event can be referred to as a topic. An event associated with a topic can be recorded on a corresponding record log. The record log can be maintained by the cluster. A producer application (e.g., an application on a client's website) can provide a record of an event to the event-driven platform, which, in turn, can add the record to the record log.

(20) When the producer provides the record of an event, a key (e.g., an application identifier or ID, such as a producer ID) is associated with the record. A consumer application can subscribe to a topic and receive the records associated with the topic. The producer and consumer applications can be referred to as producers and consumers, respectively. Collectively, producers and consumers can be referred to as streams applications. To facilitate scaling and parallelism, a topic can be divided into a set of partitions, each of which can run on a different entity of the cluster. A respective partition can maintain a subset of the event records associated with the topic. A hash of the record key can be used to index the partition. Typically, applications can send records to a particular partition selected based on the hash function.

(21) In addition, a respective application can determine a state from the records (e.g., a count or an aggregation of the event records) generated by the application. Consequently, such an application can be considered “stateful” due to the maintenance of the producer's state and can be referred to as a stateful application. A stateful application may operate without streaming events to the event-driven platform. To facilitate a backup of the state, the event-driven platform can maintain a stateful backup topic, which can be internal to the event-driven platform. The event-driven platform can then divide the stateful backup topic to generate and maintain respective backup partitions corresponding to the set of partitions. If the records generated by the producer are stored in the n th partition, the state generated by the producer is maintained in the n th backup partition. Therefore, the n th backup partition can maintain a backup state of the producer that can send event records to the n th partition.

(22) For example, if the event-driven platform is based on Kafka, the state can be a StateStore generated by the stateful applications. The stateful backup topic can then be a ChangeLog topic internally maintained by Kafka. The corresponding partition of the ChangeLog topic can maintain a backup of a StateStore of the application. If the number of partitions (or partition count) of a topic changes, the topic can be repartitioned to generate a set of new partitions. Accordingly, the event-driven platform can dynamically create a new set of backup partitions corresponding to the new partitions.

(23) The aspects described herein address the problem of dynamically scaling a topic configured on the event-driven platform by (i) redistributing the backup states among the upscaled new backup partitions such that the new backup partition corresponding to a stateful application can store a backup state of the; and (ii) transferring the record log and backup state of a stateful application to corresponding downscaled new partitions via temporary topics. For both upscaling and downscaling, a partition identifier of the new partition and new backup partition is determined based on the hash function for a respective state (e.g., based on its corresponding stateful application) of the topic. As a result, for upscaling, if a stateful application is to add records to the n th new partition, the backup state associated with the stateful application is placed in the n th new backup partition. For downscaling, if a stateful application is to add records to the n th new partition, the record log of the stateful application is placed in the n th new partition via the n th temporary partition of a temporary topic. Similarly, the stateful backup state of the stateful application is placed in the n th new backup partition via the n th temporary stateful backup partition of a temporary stateful backup topic.

(24) With existing technologies, an event-driven platform, such as Kafka, can divide a topic into a plurality of partitions. In particular, the volume of data from the stateful applications can vary, which may lead to adjustment of resource allocation to corresponding topics. For example, the stateful applications associated with the topic can continue to add event records to the topic in real time. Consequently, the volume of data associated with the topic may continue to increase. The partition count (e.g., the number of partitions) associated with the topic can determine the number of stateful applications (i.e., consumers) that may access (e.g., either read or write) the topic. A higher number of stateful applications can enable parallel data processing, which may improve the scalability of the topic. Therefore, to allow the event-driven platform to scale up to enhance the throughput, the number of partitions may be dynamically increased.

(25) On the other hand, the event records from the stateful applications associated with the topic may also be reduced. For example, some stateful applications may cease to operate, which can reduce the rate of event records arriving at the topic. The rate of event records generated by a stateful application may also be reduced. Consequently, the volume of data associated with the topic may decrease at a lower rate than anticipated. As a result, the level of parallelism facilitated by the partitions associated with the topic can be overprovisioned. Therefore, to ensure efficient utilization of the resources allocated to the topic, the number of partitions may be dynamically scaled down (or decreased).

(26) A hash function can be used to determine which stateful application may add records to which partition. The hash function can be applied to a key, such as an application ID (e.g., a producer ID), to determine the index of the partition. If the number of partitions of the topic is changed, the hash function can result in a new partition for a stateful application. For example, if the topic is divided into four partitions, the hash function can select an index from four values. If the number of partitions is upscaled to six, the hash function can select an index from six values. On the other hand, if the number of partitions is downscaled to three, the hash function can select an index from three values. As a result, when a topic is scaled up or down, the hash function may allocate a new partition to the stateful application. The index for the new partition can be different from the previous partition.

(27) If the topic is upscaled, the event-driven platform may increase the number of partitions of the topic. When the upscaling is complete, the stateful application may start sending records to the new partition. The event-driven platform can then transfer the record log from the previous partition to the new partition. In this way, the event-driven platform can dynamically redistribute the record logs so that the hash function can match the stateful application with the partition storing the corresponding record.

(28) Furthermore, during the repartitioning, the stateful applications associated with the topic can be suspended. When the stateful applications are resumed, a respective stateful application can rely on the corresponding backup partition as the source of the state because the backup partitions can maintain backup copies of the states (i.e., backup states). However, the newly created backup partitions may be empty. Therefore, the backup state associated with the stateful application may not be available at the new corresponding backup partition. Consequently, the record logs and the corresponding state are not stored based on the same index. Therefore, the existing state may become unavailable to the stateful applications. The event-driven platform may need to perform extensive computations to reconstruct the state, which can be inefficient and time intensive.

(29) Moreover, if the topic is to be downscaled, some event-driven platforms, such as Kafka, may not support the downscaling operation. As a result, a configuration command supported by the event-driven platform may not be available to an administrator to perform a downscaling operation on a topic. Therefore, to reduce the number of partitions of a topic, a new topic can typically be defined with the reduced number of partitions. However, the record logs maintained in the original topic and states maintained in the corresponding stateful backup topic are to be manually transferred to the new topic and the stateful new backup topic by an administrator. Furthermore, the stateful applications associated with the topic can already be in communication with the original topic. Hence, the stateful applications need to be reconfigured with the new topic. Transferring data between topics and configuring the stateful applications can be tedious, inefficient, and error prone.

(30) To address this issue, a partition management system operating in conjunction with the event-driven platform can dynamically scale a topic. During operation, the system can suspend the stateful applications (e.g., stateful consumer applications) associated with the topic and read from the old backup partitions. Since the stateful applications associated with a topic can write to or read from the partitions, the suspension allows the system to scale the topic (e.g., increase or decrease the number of partitions of the topic). To upscale the topic, the system may configure the event-driven platform to change the number of partitions. The configuration can be executed based on a command provided to the event-driven platform. The command can include the name or identifier of the topic and the new number of partitions.

(31) Based on the configuration, the event-driven platform can increase the number of partitions. To match the new number of partitions, the event-driven platform can internally increase the number of backup partitions to store corresponding backup states. The system can then redistribute the backup states so that the hash function can match the stateful applications with the new backup partition storing the backup state. Accordingly, the system can generate a partition identifier of the new backup partition for a respective state. A respective partition identifier can correspond to an

index generated by the hash function associated with the stateful applications. For example, if the number of partitions is increased from three to six, the partition identifiers can be from zero to five. The hash function applied to determine the new backup partition for a stateful application can be the same hash function applied to determine the new partition.

(32) Upon upscaling the topic, the system can resume the stateful applications associated with the topic. When the hash function is applied to a resumed stateful application, the corresponding index can indicate the new partition to which the stateful application is going to send subsequent records. The same index can further indicate the new backup partition from which the stateful application is going to retrieve the backup state. In addition, to ensure consistency, the same hash function can be used before and after the repartitioning. In other words, prior to the repartitioning, the same hash function might determine the previous partition allocated to the stateful application.

(33) For a stateful application and its state, if the partition identifiers of the old and new backup partitions are different, the state is stored in the new backup partition. This ensures that if the stateful application is to add records to the n th new partition, the old state associated with the stateful application is placed in the n th new backup partition. When the transfer is complete, the system can resume the operations of the stateful applications. When the stateful application adds new records to the n th partition, the event-driven platform can obtain the record log associated with the stateful application and place it in the n th partition. Hence, the new records can be incorporated into the record log.

(34) Moreover, the stateful application may update the state based on the new records. To do so, the stateful application can identify the new backup partition associated with the stateful application based on the hash function. Since the backup state has been transferred to the new backup partition, the stateful application can discover the state associated with the stateful application and retrieve the state. The stateful application can then incorporate the new records into the state. The updated state can then be automatically backed up by the event-driven platform in the new backup partition. In this way, even after the topic is dynamically upscaled, a stateful application can obtain the backup state and continue to update the state with new records.

(35) On the other hand, to downscale the topic, the system can determine the reduced number of partitions of the topic and create a corresponding temporary topic and a temporary stateful backup topic with the reduced number of partitions. For example, if the number of partitions of the topic is four and is to be reduced to three, the temporary topics are created with three partitions. For a respective stateful application associated with the topic, the system can apply the hash function to the key associated with a state of the stateful application for determining an index. Based on the index, the system can identify a temporary partition of the temporary topic and a corresponding temporary backup partition of the temporary stateful backup topic. The system can then transfer the record log of the stateful application from the current partition of the topic to the temporary partition. The system can also transfer the state associated with the stateful application from the current backup partition of the stateful backup topic to the temporary backup partition.

(36) Based on the transfer, the system can maintain respective copies of record logs and states in the temporary topics. Since the copies of record logs and states are preserved at the temporary topics, the system can delete or remove the topic and the corresponding stateful backup topic from the event-driven platform. For example, if the event-driven platform is based on Kafka, the system can remove the topic and the corresponding ChangeLog topic from Kafka. Deleting the topic and stateful backup topic can free the topic identifier, such as the topic name, at the event-driven platform. The system can then generate (or create) the topic and the corresponding stateful backup topic with the reduced number of partitions. The downscaled topics can have the same topic identifier (or name). The system can then determine respective partition identifiers for the stateful applications.

(37) Because the temporary topic and the downscaled topic have the same number of partitions, they have a one-to-one correspondence. Accordingly, the system can transfer the record logs from a

temporary partition to a corresponding new partition of the downscaled topic. In the same way, the system can transfer the states from a temporary backup partition to a corresponding new backup partition of the regenerated stateful backup topic. In other words, the system can transfer from the *n*th temporary partition to the *n*th new partition and from the *n*th temporary backup partition to the *n*th new backup partition. In this way, the system can regenerate the topics with the reduced number of partitions and populate the partitions in such a way that the record logs and states are readily available to the stateful applications.

(38) Upon downscaling the topic, the system can resume the stateful applications associated with the topic. When a resumed stateful application adds subsequent records to the new partition of the downscaled topic, the new records are added to the record log in the new partition. In the same way, since the backup state has been transferred to the new backup partition, the stateful application can discover the state associated with the stateful application and retrieve the state. The stateful application can then incorporate the new records into the state. The updated state can then be automatically backed up by the event-driven platform in the new backup partition.

(39) In this disclosure, the term “message” refers to a group of bits that can be transported together across a network. “message” should not be interpreted as limiting examples of the present invention to a particular layer of a network protocol stack. “Message” can be replaced by other terminologies referring to a group of bits, such as “packet,” “frame,” “cell,” “datagram,” or “transaction.” Furthermore, the term “port” can refer to the port that can receive or transmit data. “Port” can also refer to the hardware, software, and/or firmware logic that can facilitate the operations of that port.

(40) FIG. 1A illustrates an example of dynamic upscaling of a topic in an event-driven platform, in accordance with an aspect of the present application. An event-driven platform **100** can dynamically incorporate records of events from a number of stateful applications **112**, **114**, and **116**. These applications may or may not be streams applications operating in conjunction with event-driven platform **100**. An event can be a computer-detectable action, such as a click on a website or a customer purchasing an item. A respective stateful application can be an application that can send a message comprising respective records of one or more events to event-driven platform **100**. Event-driven platform **100** can be distributed on cluster **102**. Cluster **102** can include a plurality of computing devices, such as servers, and storage devices, such as persistent storage arrays.

Typically, upon receiving a message, event-driven platform **100** can add the record of an event (or event record) indicated in the message to an event log stored in a persistent storage of cluster **102**.

(41) A respective type of event being stored at event-driven platform **100** can be referred to as a topic. Suppose that a topic **130** is defined in event-driven platform **100**. Stateful applications **112**, **114**, and **116** can be associated with topic **130**. Therefore, a respective event associated with topic **130** from stateful applications **112**, **114**, and **116** can be recorded on a record log of topic **130**. The record log can be maintained by event-driven platform **100** in cluster **102**. The resources allocated for storing topic **130** can be managed by event-driven platform **100**. To facilitate scaling and parallelism, topic **130** can be divided into a set of partitions **131**, **132**, and **133**, each of which can run on a different entity of cluster **102**. The event records from a particular stateful application are maintained at one of partitions **131**, **132**, and **133**. Therefore, a respective partition may maintain respective event logs of an exclusive set of stateful applications.

(42) For example, when stateful application **112** sends a message comprising the record of an event to event-driven platform **100**, a key (e.g., an identifier of stateful application **112**) is associated with the record. A hash of the key can be used to generate an index that can select one of partitions **131**, **132**, and **133**. In this example, partitions **131**, **132**, and **133** can be selected for storing record logs of stateful applications **112**, **114**, and **116**, respectively, based on the hashing. A hash function can be used to determine the index. The hash function can use the key and the number of partitions as inputs and generate an index as an output. If the number of partitions is *N*, the index value can fall within a range of $[0, N]$. A consumer application can subscribe to topic **130** and receive the records

associated with topic **130**.

(43) Furthermore, stateful applications **112**, **114**, and **116** can determine states **122**, **124**, and **126**, respectively, from their records. For example, if stateful application **112** calculates the sales volume of an online shop, the event records can indicate revenue generated from individual sales, and state **122** can then be the total revenue generated. Consequently, stateful applications **112**, **114**, and **116** can be considered as “stateful applications” due to the maintenance of states **122**, **124**, and **126**, respectively. To facilitate a backup of states **122**, **124**, and **126**, event-driven platform **100** can maintain a stateful backup topic **140**, which can be internal to event-driven platform **100**. To maintain consistency with topic **130**, event-driven platform **130** can divide stateful backup topic **140** to generate backup partitions **141**, **142**, and **143**.

(44) Backup partitions **141**, **142**, and **143** can correspond to partitions **131**, **132**, and **133**, respectively. Therefore, for each of stateful applications **112**, **114**, and **116**, the same index can be used to select one of partitions **131**, **132**, and **133** and one of backup partitions **141**, **142**, and **143**. Accordingly, since the event records generated by stateful application **112** can be stored in partition **131**, a backup copy of state **122** generated by stateful application **112** can be maintained in backup partition **141**. In the same way, respective backup copies of states **124** and **126** can be maintained in backup partitions **142** and **143**, respectively. If stateful application **112** becomes unavailable (e.g., due to an error or an upgrade), stateful application **112** can retrieve the backup copy of state **122** from backup partition **141**.

(45) Here, if event-driven platform **100** is based on Kafka, states **122**, **124**, and **126** can be StateStores generated by stateful applications **112**, **114**, and **116**, respectively. Stateful backup topic **140** can then be a ChangeLog topic internally maintained by Kafka. Partitions **141**, **142**, and **143** of the ChangeLog topic can maintain a backup copy of StateStores **122**, **124**, and **126**, respectively. If the number of partitions of topic **130** changes due to upscaling or downscaling, topic **130** can be repartitioned to generate a set of new partitions. Accordingly, event-driven platform **100** can dynamically create a new set of backup partitions of stateful backup topic **140** corresponding to the new partitions.

(46) With existing technologies, stateful applications **112**, **114**, and **116** can continue to add event records to topic **130** in real time. Consequently, the volume of data associated with topic **130** may continue to increase. Because topic **130** has been divided into three partitions (i.e., partitions **131**, **132**, and **133**), there can be three stateful applications (i.e., consumers) that can access topic **130**. A higher number of stateful applications can enable parallel data processing, which may improve the scalability for topic **130**. Therefore, to allow event-driven platform **100** to scale up to enhance the throughput, the number of partitions of topic **130** may be dynamically increased.

(47) On the other hand, the event records from stateful applications **112**, **114**, and **116** may also be reduced. For example, one of stateful applications **112**, **114**, and **116** can stop operations. Consequently, the rate of event records arriving at topic **130** can decrease as well. Furthermore, the rate of event records generated by a stateful application, such as stateful application **116**, may also be reduced, and the volume of data associated with topic **130** may decrease. As a result, the level of parallelism facilitated by partitions **131**, **132**, and **134** can be overprovisioned. Therefore, to ensure efficient utilization of the resources of cluster **102** allocated to topic **130**, the number of partitions may be dynamically scaled down (or decreased).

(48) If the topic is upscaled, event-driven platform **100** may increase the number of partitions of topic **130**. Suppose that the number of partitions of topic **130** is increased to six. An administrator may issue a configuration command to event-driven platform **100** to upscale topic **130** to six. The command can be supported by event-driven platform **100** and include the name or identifier of topic **100** and the new number of partitions (e.g., six) as inputs. Accordingly, event-driven platform **100** can then generate partitions **131**, **132**, **133**, **134**, **135**, and **136** for topic **130**. To match the new number of partitions, event-driven platform **100** can internally upscale stateful backup topic **140** and generate backup partitions **141**, **142**, **143**, **144**, **145**, and **146** for stateful backup topic **140**.

Before the upscaling, the hash function can produce an index from three values. When the number of partitions is upscaled to six, the hash function can select an index from six values. As a result, when topic **130** is upscaled, the hash function may allocate a new partition to a stateful application. The index for the new partition can be different from the previous partition.

(49) For example, before the upscaling, when the hash function is applied for stateful application **112**, the index can correspond to partition **131**. After the upscaling, the index may still correspond to partition **131**. On the other hand, before the upscaling, when the hash function is applied for stateful application **114**, the index can correspond to partition **132**. Accordingly, when the upscaling is complete, stateful application **114** may start sending records to partition **133**. Event-driven platform **100** can then transfer the record log from partition **131** to partition **133**. In this way, event-driven platform **100** can dynamically redistribute the record logs so that the hash function can match the stateful application with the partition storing the corresponding record.

(50) Furthermore, during the upscaling of topic **130**, stateful applications **112**, **114**, and **116** can be suspended. When stateful applications **112**, **114**, and **116** are resumed, a respective stateful application can rely on the corresponding backup partition as the source of states **122**, **124**, and **126**. However, newly created backup partitions **141**, **142**, **143**, **144**, **145**, and **146** may be empty. Therefore, the backup copies of states **122**, **124**, and **126** may not be available at the new backup partitions. Consequently, states **122**, **124**, and **126** may not be readily available to stateful applications **112**, **114**, and **116**, respectively. Event-driven platform **100** may need to perform extensive computations to reconstruct states **122**, **124**, and **126**, which can be inefficient and time intensive.

(51) To address this issue, a partition management system **110** operating in conjunction with event-driven platform **100** can dynamically upscale topic **130**. Respective instances of system **110** can operate with stateful applications **112**, **114**, and **116** to facilitate the efficient scaling of topic **130**. In some examples, the instances of system **110** may be collocated with stateful applications **112**, **114**, and **116** (e.g., operate on the same computing system) and may operate as a root user on these applications. During operation, the system can suspend stateful applications **112**, **114**, and **116** and read from partitions **141**, **142**, and **143**. When event-driven platform **100** increases the number of partitions of topic **130**, event-driven platform **100** can internally increase the number of backup partitions of stateful backup topic **140**. System **110** can then redistribute the backup copies of states **122**, **124**, and **126** so that the hash function can match them with stateful applications **112**, **114**, and **116**.

(52) Accordingly, system **110** can generate respective partition identifiers (or partition numbers) for states **122**, **124**, and **126** (e.g., based on the key associated with stateful application **112**, **114**, and **116**, respectively). The partition identifier can identify backup partitions from partitions **141**, **142**, **143**, **144**, **145**, and **146**. A respective partition identifier can correspond to an index generated by the hash function. Since the number of partitions has increased to six, the partition identifiers can be from zero to five. The partition identifiers can correspond to the indices generated by the hash function. If the indices generated for stateful application **114** are zero and one, partitions **141** and **142** are selected for stateful application **112**. Therefore, when the hash function is applied to stateful application **112**, the corresponding indices zero and one can identify partitions **131** and **132**, to which stateful application **112** is going to send subsequent event records. The same indices can further indicate partitions **141** and **142**, from which stateful application **112** is going to retrieve the backup of state **122**. In addition, to ensure consistency, the same hash function can be used before and after the scaling.

(53) In the same way, for stateful application **114**, the hash function can select partitions **133** and **134** of upscaled topic **130** and partitions **143** and **144** of upscaled stateful backup topic **140**. Here, the partition identifiers of old backup partition **142** and new backup partitions **143** and **144** for state **124** can be different. Therefore, system **110** can transfer state **124** from old backup partition **142** to new backup partitions **143** and **144**. Moreover, for stateful application **116**, the hash function can

select partitions **135** and **136** of upscaled topic **130** and partitions **145** and **146** of upscaled stateful backup topic **140**. Here, the partition identifiers of old backup partition **143** and new backup partitions **145** and **146** for state **126** can be different. Therefore, system **110** can transfer state **126** from old backup partition **143** to new backup partitions **145** and **146**.

(54) When the transfer of states **124** and **126** is complete, stateful applications **112**, **114**, and **116** can be resumed. If stateful application **114** updates state **124** based on new records, stateful application **114** can identify new backup partition **143** based on the hash function. Since the backup copy of state **124** has been transferred to new backup partition **143**, stateful application **114** can discover state **124** and retrieve it. Stateful application **114** can then incorporate the new records into state **124**. The updated state can then be automatically backed up by event-driven platform **100** in new backup partitions **143** and **144**. In this way, even after topic **130** is dynamically upscaled, stateful application **114** can obtain state **124** and continue to update state **124** with new records.

(55) FIG. 1B illustrates an example of dynamic downscaling of a topic in an event-driven platform, in accordance with an aspect of the present application. If the topic is to be downscaled, event-driven platform **100** may not support a downscaling operation. In other words, event-driven platform **100** may not be equipped with a command that can downscale topic **130**. Therefore, to reduce the number of partitions of topic **130**, a new topic can typically be defined with the reduced number of partitions. However, the record logs maintained in topic **130** and copies of states **122**, **124**, and **126** maintained in stateful backup topic **140** are to be manually transferred to the new topic and the stateful new backup topic by an administrator. Furthermore, stateful applications **112**, **114**, and **116** can already be in communication with topic **130**. Hence, stateful applications **112**, **114**, and **116** need to be reconfigured with the new topic. Transferring data between topics and configuring stateful applications **112**, **114**, and **116** can be tedious, inefficient, and error prone.

(56) To address this issue, to downscale topic **130**, system **110** can determine the reduced number of partitions. System **110** can then create a corresponding temporary topic **150** and a temporary stateful backup topic **160** with two partitions. In this example, the number of partitions of topic **130** can be three and is to be reduced to two. Therefore, temporary topic **150** can be created with two partitions **151** and **152**, and temporary stateful backup topic **160** can be created with two partitions **161** and **162**. The system can then apply the hash function to the respective application IDs of stateful applications **112**, **114**, and **116** and determine corresponding indices. Suppose that the indices for stateful applications **112**, **114**, and **116** are zero, one, and one, respectively. Accordingly, system **110** can select temporary partition **151** and temporary backup partition **161** for stateful application **112**, and temporary partition **152** and temporary backup partition **162** for stateful applications **114** and **116**.

(57) System **110** can then transfer the record log of stateful application **112** from partition **141** to temporary partition **151**, and the copy of state **122** from backup partition **141** to temporary backup partition **161**. Similarly, system **110** can transfer the record log of stateful application **114** from partition **142** to temporary partition **152**, and the copy of state **124** from backup partition **142** to temporary backup partition **162**. System **110** can also transfer the record log of stateful application **116** from partition **143** to temporary partition **152**, and the copy of state **126** from backup partition **143** to temporary backup partition **162**. Based on the transfer, system **110** can preserve respective copies of record logs and states in temporary topics **150** and **160**.

(58) System **110** can then delete or remove topic **130** and stateful backup topic **140** from event-driven platform **100**. For example, if event-driven platform **100** is based on Kafka, system **110** can remove topic **130** and ChangeLog topic **140** from Kafka. Deleting topic **130** and stateful backup topic **140** can free the topic identifier, such as the topic name, at event-driven platform **100**. System **110** can generate topic **130** and stateful backup topic **140** with two partitions. In this way, system **110** can regenerate downscaled topic **130** and downscaled stateful backup topic **140**, where they are scaled down from three partitions to two partitions. Here, downscaled topic **130** can have the same topic identifier (or name) as the original topic **130**. For a respective stateful application, system **110**

can determine partition identifiers of the partitions of topics **130** and **140**, which can correspond to the indices generated by the hash function. Here, the partition identifiers for partitions **141** and **142** in downscaled backup partition **140** can be zero and one.

(59) Because temporary topic **150** and downscaled topic **130** have the same number of partitions, the partitions of topics **130** and **150** have a one-to-one correspondence. Accordingly, system **110** can transfer record logs from temporary partitions **151** and **152** to partitions **131** and **132** of downscaled topic **130**, respectively. In the same way, system **110** can transfer the states from temporary backup partitions **161** and **162** to backup partitions **141** and **142** of downscaled stateful backup topic **140**. In this way, system **110** can regenerate topics **130** and **140** with the reduced number of partitions and populate their respective partitions in such a way that the record logs and states are available to stateful applications **112**, **114**, and **116**.

(60) Upon downscaling topic **130**, system **110** can resume stateful applications **112**, **114**, and **116**. Since downscaled topic **130** can be generated based on the same topic identifier or topic name on event-driven platform **100**, stateful applications **112**, **114**, and **116** can readily access downscaled topic **130** without requiring application-level reconfiguration. Hence, when resumed stateful application **116** adds subsequent records to partition **132** of downscaled topic **130**, the new records are added to the record log of stateful application **116** in partition **132**. In the same way, since the backup copy of state **126** has been transferred to backup partition **142**, stateful application **116** can discover state **126** and retrieve it. Stateful application **116** can then incorporate the new records into state **126**. The updated state **126** can then be automatically backed up by event-driven platform **100** in partition **142** of downscaled stateful backup topic **140**.

(61) FIG. 2A illustrates an example of a partition management system increasing the number of partitions of a topic in an event-driven platform, in accordance with an aspect of the present application. A topic **214** can be defined on an event-driven platform **200**. A stateful application **212** can be associated with topic **214**. Stateful application **212** can be a producer that sends a set of messages to topic **214** to store a stream of event records. Stateful application **212** can be a stateful application that can maintain a state (e.g., a Kafka StateStore) determined based on the event records. To maintain a backup copy of the state, event-driven platform **200** can internally generate a stateful backup topic **216**.

(62) A partition management system **230** can operate in conjunction with event-driven platform **200** to facilitate upscaling and downscaling of topic **214**. System **230** can operate with stateful application **212** to facilitate the efficient scaling of topic **214**. In some examples, system **230** may be collocated with stateful application **212** and may operate as a root user of stateful application **212**. System **230** can include a partition management proxy **202**, an application controller **204**, and a topic manager **206**.

(63) Proxy **202** can facilitate an interface between a user (e.g., an administrator managing topic **214**) and system **230**. Proxy **202** may invoke a respective other component of system **230** used for scaling topic **214**. The user can provide a command to upscale topic **214** via proxy **202**. The command can include the topic identifier or name of topic **214** and the increased number of partitions for topic **214**. Proxy **202** may present a user interface to the user to obtain the command. Examples of the user interface can include, but are not limited to, a textual interface, a graphical user interface (GUI), a movement-based interface, a gesture-based interface, a virtual or augmented reality interface, and a touch interface.

(64) Application controller **204** can suspend and resume a respective stateful application associated with topic **214**. Furthermore, topic manager **206** can manage the repartitioning (or altering) topic **214** and stateful backup topic **216**. Topic manager **206** can interact with event-driven platform **200** using an application programming interface (API) supported by event-driven platform **200**. For example, if event-driven platform **200** is based on Kafka, the API can include Kafka administrative APIs. During operation, upon receiving the command from the user, proxy **202** can invoke application controller **204** (operation **232**). The invocation may include an instruction to suspend

the stateful applications associated with topic **214**.

(65) Accordingly, application controller **204** can suspend stateful application **212** (operation **234**). The suspension can cause stateful application **212** to stop sending event records to topic **214**. Application controller **204** may issue a command facilitated by event-driven platform **200** to suspend stateful application **212**. In response, application controller **204** can receive an indication (e.g., a return value of the command) from stateful application **212** indicating the suspension (operation **236**). The suspension of stateful application **212** can complete the operations associated with the invocation of application controller **204**, which can then return the execution back to proxy **202** (operation **238**).

(66) Subsequently, proxy **202** can invoke topic manager **206** (operation **240**). The invocation may include an instruction to upscale topic **214** to the new number of partitions, as indicated in the command from the user. Topic manager **206** can then use the APIs to reconfigure topic **214** on event-driven platform **200** to increase the partition count (operation **242**). Based on the reconfiguration, event-driven platform **200** can increase the number of partitions of topic **214**. This upscaling of topic **214** may be managed and handled by event-driven platform **200**. When event-driven platform **200** completes the upscaling process, event-driven platform **200** can indicate, via the APIs, that the process is complete, and the partition count is increased (operation **244**).

(67) Similarly, topic manager **206** can use the APIs to reconfigure stateful backup topic **216** on event-driven platform **200** to increase the partition count to match it with the partition count of topic **214** (operation **246**). The upscaling of back topic **216** may also be managed and handled by event-driven platform **200**. Upon completion, event-driven platform **200** can indicate, via the APIs, that the upscaling process for stateful backup topic **216** is complete, and the partition count is increased (operation **248**). The upscaling of topic **214** and stateful backup topic **216** can complete the operations associated with the invocation of topic manager **206**, which can then return the execution back to proxy **202** (operation **250**).

(68) System **230** can then ensure that the new backup partition of upscaled stateful backup topic **216** allocated to stateful application **212** stores the state of stateful application **212** by redistributing the states (operation **252**). The redistribution of the states is further described in conjunction with FIG. 2B. After the redistribution, the state can be accessible to stateful application **212** from the corresponding new backup partition of upscaled stateful backup topic **216**. Therefore, stateful application **212** can be ready to resume operation. Proxy **202** can then invoke application controller **204** (operation **254**). The invocation may include an instruction to resume the stateful applications associated with topic **214**.

(69) Accordingly, application controller **204** can resume stateful application **212** (operation **256**). The resumption can cause stateful application **212** to start sending event records to topic **214**. Application controller **204** may issue a command to start stateful application **212**. In response, application controller **204** can receive an indication from stateful application **212** indicating the resumption (operation **258**). The resumption of stateful application **212** can complete the operations associated with the invocation of application controller **204**, which can then return the execution back to proxy **202** (operation **260**).

(70) FIG. 2B illustrates an example of redistributing backup states among the increased number of partitions of a topic in an event-driven platform, in accordance with an aspect of the present application. System **230** can further include a record publisher **208** and a partition manager **210**. Here, partition manager **210** can read states from the backup partitions of stateful backup topic **216**. To do so, partition manager **210** can incorporate an instance of a partitioner **218** of event-driven platform **200**. Partitioner **218** can determine which stateful application is allocated to which partition. Partitioner **218** can implement the hash function that can be applied to associate a stateful application, its record logs, and its state with a corresponding partition and backup partition. Therefore, partition manager **210** can be used to read from stateful backup topic **216**. Record publisher **208** can “publish” a record log or state to a partition. The publishing process can include

transferring the record log or state to the partition (e.g., from a previous partition).

(71) To redistribute a state, proxy **202** can invoke partition manager **210** (operation **262**). The invocation may include an instruction to redistribute the states of stateful backup topic **216**. Partition manager **210** can then use partitioner **218** to read a state from a previous backup partition of stateful backup topic **216** (operation **264**). It should be noted that the new backup partitions of upscaled stateful backup topic **216** can be empty. Because partitioner **218** can support the hash function, partition manager **210** can also use the hash function to determine a new partition identifier allocated for the state (operation **266**).

(72) Based on the new partition identifier, partition manager **210** may determine that the partition identifiers of the previous and new backup partitions are different (operation **268**). The difference of the identifiers can indicate that, upon resumption, the corresponding stateful application is to seek the backup copy of its state from a different backup partition. Therefore, partition manager **210** can invoke record publisher **208** (operation **270**). The invocation may include an instruction to publish the state to the new backup partition corresponding to the new partition identifier. Based on the invocation, record publisher **208** can transfer the state to the new backup partition of upscaled stateful backup topic **216** (operation **272**).

(73) Based on the new partition identifier, partition manager **210** may also determine that the partition identifiers of the previous and new backup partitions are the same (operation **274**). Consequently, upon resumption, the corresponding stateful application is to seek the backup copy of its state from the same backup partition. Therefore, partition manager **210** may skip the transfer of the state (operation **276**). Partition manager **210** can repeat the process of generating a partition identifier and publishing the state based on the partition identifier for each state in stateful backup topic **216** (operation **278**). When the process is repeated for each state, the redistribution of states can be complete. Partition manager **210** can then return the execution back to proxy **202** (operation **280**).

(74) FIG. 3A illustrates an example of a partition management system generating temporary topics for decreasing the number of partitions of a topic in an event-driven platform, in accordance with an aspect of the present application. A topic **314** can be defined on an event-driven platform **300**. A stateful application **312** can be associated with topic **314**. Stateful application **312** can be a producer that sends messages to topic **314** for storing event records. Stateful application **312** can be a stateful application that can maintain a state (e.g., a Kafka StateStore) determined based on the event records. To maintain a backup copy of the state, event-driven platform **300** can internally generate a stateful backup topic **316**.

(75) A partition management system **330** can operate in conjunction with event-driven platform **300** to facilitate upscaling and downscaling of topic **314**. System **330** can operate with stateful application **312** to facilitate the efficient scaling of topic **314**. In some examples, system **330** may be collocated with stateful application **312** and may operate as a root user of stateful application **312**. System **330** can include a partition management proxy **302**, an application controller **304**, and a topic manager **306**. Proxy **302** can facilitate an interface between a user (e.g., an administrator managing topic **214**) and system **330**. Proxy **302** may invoke a respective other component of system **330** used for scaling topic **314**. The user can provide a command to upscale topic **314** via proxy **302**. The command can include the topic identifier or name of topic **314** and the reduced number of partitions for topic **314**. Proxy **302** may present a user interface to the user to obtain the command.

(76) Application controller **304** can suspend and resume a respective stateful application associated with topic **314**. Furthermore, topic manager **306** can manage the repartitioning (or altering) topic **314** and stateful backup topic **316**. Topic manager **306** can interact with event-driven platform **300** using an API supported by event-driven platform **300**. For example, if event-driven platform **300** is based on Kafka, the API can include Kafka administrative APIs. During operation, upon receiving the command from the user, proxy **302** can invoke application controller **304** (operation **331**). The

invocation may include an instruction to suspend the stateful applications associated with topic **314**.

(77) Accordingly, application controller **304** can suspend stateful application **312** (operation **332**). The suspension can cause stateful application **312** to stop sending event records to topic **314**. Application controller **304** may issue a command facilitated by event-driven platform **300** to suspend stateful application **312**. In response, application controller **304** can receive an indication (e.g., a return value of the command) from stateful application **312** indicating the suspension (operation **333**). The suspension of stateful application **312** can complete the operations associated with the invocation of application controller **304**, which can then return the execution back to proxy **302** (operation **334**).

(78) Subsequently, proxy **302** can invoke topic manager **306** (operation **336**). The invocation may include an instruction to generate downscaled stateful backup topics for topic **314** and stateful backup topic **316**. Topic manager **306** can then use the APIs to generate a temporary stateful backup topic **320** on event-driven platform **300** with the reduced number of partitions (i.e., with the decreased partition count) (operation **336**). Temporary stateful backup topic **320** can operate as a placeholder for the states in stateful backup topic **316** during the downscaling of topic **314**. Topic manager **306** may issue a command facilitated by the APIs for generating temporary stateful backup topic **320**. Based on the command, event-driven platform **300** can generate temporary stateful backup topic **320**. When event-driven platform **300** completes the generation process, event-driven platform **300** can indicate, via the APIs, that the process is complete, and temporary stateful backup topic **320** is generated (operation **337**).

(79) Furthermore, topic manager **306** can use the APIs to generate a temporary topic **318** on event-driven platform **300** with the reduced number of partitions (operation **338**). Temporary topic **318** can operate as a placeholder for the record logs in topic **314** during its downscaling. Topic manager **306** may issue a command facilitated by the APIs for generating temporary topic **318**. Based on the command, event-driven platform **300** can generate temporary topic **318**. When event-driven platform **300** completes the generation process, event-driven platform **300** can indicate, via the APIs, that the process is complete, and temporary topic **318** is generated (operation **339**). When the generation of temporary topic **318** and temporary stateful backup topic **320** are complete, topic manager **306** can return the execution back to proxy **302** (operation **340**).

(80) FIG. 3B illustrates an example of a partition management system transferring record logs and states to temporary topics in an event-driven platform, in accordance with an aspect of the present application. System **330** can further include a record publisher **308** and a partition manager **310**. Here, partition manager **310** can read record logs from the partitions of topic **314** and states from the backup partitions of stateful backup topic **316**. To do so, partition manager **310** can incorporate an instance of a partitioner of event-driven platform **300**. The partitioner can implement the hash function that can be applied to associate a stateful application, its record logs, and its state with a corresponding partition and backup partition. Record publisher **308** can publish a record log or state to a partition. The publishing process can include transferring the record log or state to a partition.

(81) To transfer data from topics **314** and **316** to temporary topics **318** and **320**, respectively, proxy **302** can invoke partition manager **310** (operation **341**). The invocation may include an instruction to transfer record logs of topic **314** to temporary topic **318** and states of stateful backup topic **316** to temporary stateful backup topic **320**. Partition manager **310** can then read record logs from topic **316** (operation **342**). To do so, partition manager **310** can read from a respective partition of topic **316**. Based on the reading, partition manager **310** can obtain the record logs (operation **343**). In the same way, partition manager **310** can read the backup copies of the states from the backup partitions of stateful backup topic **316** (operation **344**). Based on the reading, partition manager **310** can obtain the backup copies of the states (operation **345**).

(82) Upon obtaining the record logs and states from topics **314** and **316**, respectively, partition

manager **310** is to place the record logs and states in temporary topics **318** and **320**, respectively. Therefore, partition manager **310** can invoke record publisher **308** (operation **346**). The invocation may include an instruction to publish the record logs and states to the temporary partitions. Based on the invocation, record publisher **308** can transfer the record logs obtained from topic **314** to the temporary partitions of temporary topic **318** (operation **347**).

(83) For a respective record log, system **330** can apply the hash function to a key (e.g., an identifier of a corresponding stateful application) of the record log and determine an index. The index can correspond to the number of partitions of temporary topic **318**. Record publisher **308** can select which temporary partition a record log should be transferred to based on the index associated with the record log. When the transfer is complete, record publisher **308** can determine that all record logs are backed up in temporary topic **318** (operation **348**).

(84) Furthermore, record publisher **308** can also transfer the states obtained from temporary topic **316** and transfer them to the temporary backup partitions of temporary stateful backup topic **320** (operation **349**). For a respective state, system **330** can apply the hash function to select which temporary backup partition a state should be transferred. When the transfer is complete, record publisher **308** can determine that all states are backed up in temporary stateful backup topic **320** (operation **350**). When the record logs and states are backed up, the operations associated with the invocation of record publisher **308** are complete. Record publisher **308** can then return the execution back to partition manager **310** (operation **351**). At this point, partition manager **310** has completed the transfer of record logs from topic **314** to temporary topic **318** and states from stateful backup topic **316** to temporary stateful backup topic **320**. Accordingly, partition manager **310** can return the execution back to proxy **302** (operation **352**).

(85) FIG. 3C illustrates an example of a partition management system regenerating topics in an event-driven platform, in accordance with an aspect of the present application. To downscale topic **314**, system **330** can delete topic **314** and stateful backup topic **316** and regenerate them. Therefore, proxy **302** can invoke topic manager **306** (operation **353**). The invocation may include an instruction to regenerate downscaled topic **314** and a corresponding downscaled stateful backup topic **316**. Topic manager **306** can then use the APIs to delete topic **314** from event-driven platform **300** (operation **354**). To ensure correspondence between topic **314** and stateful backup topic **316**, topic manager **306** can also delete stateful backup topic **316** from event-driven platform **300** (operation **355**).

(86) Topic manager **306** may issue a command facilitated by the APIs for deleting topic **314**. Based on the command, event-driven platform **300** can internally delete stateful backup topic **316** and indicate that stateful backup topic **316** is deleted (operation **356**). In addition, event-driven platform **300** can also delete topic **314** based on the command and indicate that topic **314** is deleted (operation **357**). Deleting topic **314** can free the topic identifier or name of topic **314** for reuse.

Therefore, topic **314** can be regenerated on event-driven platform **300** with the same identifier.

(87) Accordingly, topic manager **306** can use the APIs to regenerate topic **314** on event-driven platform **300** with the reduced number of partitions (i.e., with the decreased partition count) (operation **358**). Topic manager **306** may issue a command facilitated by the APIs for regenerating topic **314**. Based on the command, event-driven platform **300** can generate topic **314** and indicate that topic **314** is regenerated (operation **359**). In this way, topic **314** can be regenerated with the same topic identifier and the reduced number of partitions. Furthermore, topic manager **306** can use the APIs to regenerate stateful backup topic **316** on event-driven platform **300** with the reduced number of partitions (operation **360**). When event-driven platform **300** completes the generation process, event-driven platform **300** can indicate, via the APIs, that stateful backup topic **316** is regenerated (operation **361**). When topic **314** and stateful backup topic **316** are regenerated, topic manager **306** can return the execution back to proxy **302** (operation **362**).

(88) FIG. 3D illustrates an example of a partition management system transferring record logs and states to regenerated topics in an event-driven platform, in accordance with an aspect of the present

application. To transfer data from temporary topics **318** and **320** to regenerated topics **314** and **316**, respectively, proxy **302** can invoke partition manager **310** (operation **363**). The invocation may include an instruction to transfer record logs of temporary topic **318** to regenerated topic **314** and states of temporary stateful backup topic **320** to stateful backup topic **316**. Partition manager **310** can then read record logs from temporary topic **318** (operation **364**). To do so, partition manager **310** can read from a respective temporary partition of temporary topic **318**. Based on the reading, partition manager **310** can obtain the record logs backed up in temporary topic **318** (operation **365**). (89) In the same way, partition manager **310** can read the states from the temporary backup partitions of temporary stateful backup topic **320** (operation **366**). Based on the reading, partition manager **310** can obtain the states (operation **367**). Upon obtaining the record logs and states from temporary topics **318** and **320**, respectively, partition manager **310** is to place the record logs and states in regenerated topics **314** and **316**, respectively. Therefore, partition manager **310** can invoke record publisher **308** (operation **368**). The invocation may include an instruction to publish the record logs and states to the regenerated partitions. Based on the invocation, record publisher **308** can transfer the record logs obtained from temporary topic **318** to the partitions of regenerated topic **316** (operation **369**).

(90) Because the partitions of regenerated topic **314** and temporary topic **318** have a one-to-one correspondence, the record logs can be transferred between the partitions with the same partition number. For example, the record logs from the *n*th temporary partition of temporary topic **318** can be transferred to the *n*th partition of regenerated topic **316**. When the transfer is complete, record publisher **308** can determine that all record logs are instated in regenerated topic **316** (operation **370**). When the record logs and states are reinstated, the operations associated with the invocation of record publisher **308** are complete. Record publisher **308** can then return the execution back to partition manager **310** (operation **373**). At this point, partition manager **310** has completed the transfer of record logs and states to regenerated topic **314** and stateful backup topic **316**.

Accordingly, partition manager **310** can return the execution back to proxy **302** (operation **374**).

(91) FIG. 3E illustrates an example of a partition management system removing temporary topics in an event-driven platform, in accordance with an aspect of the present application. Once the record logs and states are transferred from temporary topic **318** and temporary stateful backup topic **320**, these topics may no longer be needed. Hence, proxy **302** can invoke topic manager **306** to remove these temporary topics (operation **375**). Topic manager **306** can then use the APIs to delete temporary topic **318** from event-driven platform **300** (operation **376**). Topic manager **306** may issue a command facilitated by the APIs for removing temporary topic **318**. Upon deleting temporary topic **318**, event-driven platform **300** can indicate that temporary topic **318** is deleted (operation **377**).

(92) Furthermore, topic manager **306** can use the APIs to delete a temporary stateful backup topic **320** from event-driven platform **300** (operation **378**). Topic manager **306** may issue a command facilitated by the APIs for generating temporary stateful backup topic **320**. Upon deleting temporary stateful backup topic **320**, event-driven platform **300** can indicate that temporary stateful backup topic **320** is deleted (operation **379**). When temporary topic **318** and temporary stateful backup topic **320** are deleted, topic manager **306** can return the execution back to proxy **302** (operation **380**). In this way, system **330** can delete temporary topics **318** and **320** to free the resources of event-driven platform **300**.

(93) At this point, stateful application **312** can start sending event records to event-driven platform **300**. Therefore, proxy **302** can invoke application controller **304** (operation **381**). The invocation may include an instruction to resume the stateful applications associated with topic **314**. Application controller **304** can then resume stateful application **312** (operation **382**). The resumption can cause stateful application **312** to start sending event records to topic **314**. Application controller **304** may issue a command facilitated by event-driven platform **300** to start stateful application **312**. In response, application controller **304** can receive an indication from

stateful application **312** indicating the resumption (operation **383**). The resumption of stateful application **312** can complete the operations associated with the invocation of application controller **304**, which can then return the execution back to proxy **302** (operation **384**).

(94) FIG. **4A** presents a flowchart illustrating the process of a partition management system upscaling a topic in an event-driven platform, in accordance with an aspect of the present application. During operation, the system can suspend a stateful application associated with a topic and a stateful backup topic defined on an event-driven platform such that the topic corresponds to a stream of computer-detectable events associated with the stateful application (operation **402**). The suspension can stop further updates from the stateful application to the event logs. Therefore, the system can perform the upscaling of the topic while the stateful application is suspended. The system can then modify the configuration of the topic at the event-driven platform to allow a set of partitions associated with the topic where the modified configuration generates a new set of backup partitions provided by the event-driven platform for the stateful backup topic (operation **404**).

(95) The system can identify the state generated based on the event records from the stateful application in a previous backup partition of the stateful backup topic (operation **406**). The previous backup partition can be a partition of the backup partition before the modification of the configuration of the topic. The system can determine a partition identifier associated with the state such that the partition corresponds to the new backup partition (operation **408**). The system can apply a hash function for the state based on the new number of partitions, which can be the number of partitions after the modification of the configuration of the topic. Based on the resultant hash value, the system can determine the partition identifier. The system can identify the new backup partition allocated to the stateful application (operation **410**). The system can use the partition identifier to identify the new backup partition.

(96) The previous backup partition can also be associated with a partition identifier. The system can compare the corresponding partition identifiers of the previous and new partitions (operation **412**). The system can then determine whether the backup new partition is different than the previous backup partition (operation **414**). Here, if the partition identifiers of the previous and new partitions are different, the system can determine that the new backup partition is different. If the new backup partition is different, the stateful application has been allocated a new backup partition for its state. Therefore, the system can transfer the state from the previous backup partition to the new backup partition (operation **416**).

(97) If the new backup partition has the same partition identifier (operation **414**) or upon transferring the state to the new backup partition (operation **416**), the system can resume the stateful application for generating an updated state, which is backed up at the new backup partition, based on the state and the subsequent event records associated with the stateful application (operation **418**). The new partition can then maintain an updated event log based on the subsequent event records from the producer application. On the other hand, the stateful application can maintain the updated state, which is backed up at the new backup partition, generated based on the state and the subsequent event records. In this way, even when the state is assigned a new backup partition, the stateful application can retrieve the state after resumption.

(98) FIG. **4B** presents a flowchart illustrating the process of a stateful application updating a state based on a new backup partition in an event-driven platform, in accordance with an aspect of the present application. Here, the stateful application can be a resumed stateful application associated with an upscaled topic. During operation, the stateful application can retrieve the state from the new backup partition (operation **432**). The system can place the state in the new partition backup partition prior to the resumption of the stateful application. Therefore, the stateful application can retrieve the state from the new partition backup. The system can then update the state based on subsequent event records (operation **434**). These subsequent event records can be indicative of the event associated with the application after the resumption of operations.

(99) The stateful application can send, to the event-driven platform, respective messages

comprising the subsequent event records, which are stored in the new partition (operation **436**). The hash function associated with the stateful application can generate an index associated with the new partition. The system can then allow the event-driven platform to transfer the event records from the previous partition to the new partition upon receiving the subsequent event record (operation **438**). When the event-driven platform receives the messages with the subsequent event record, it can determine that a new partition has been allocated to the stateful application based on the hash function. The event-driven platform can then transfer the event log comprising the event records to the new partition to allow the stateful application to incorporate the subsequent event records.

(100) FIG. 4C presents a flowchart illustrating the process of a stateful application identifying a new backup partition in an event-driven platform, in accordance with an aspect of the present application. The stateful application can select the hash function used to determine the new partition for storing event records from the stateful application (operation **452**). There can be multiple hash functions supported by the event-driven platform. The stateful application can select one of the supported ones. However, the stateful application can select the same hash function for identifying a partition of a topic and a backup partition of a stateful backup topic. In this way, the same partition identifier can correspond to the partition and the backup partition.

(101) Upon selecting the hash function, the stateful application can determine a key associated with the stateful application (operation **454**). The key can be an identifier of the stateful application. The stateful application can then apply the hash function on the key associated with the stateful application to generate a hash value (operation **456**). The hash value can be an index based on which a partition can be selected. Therefore, the index can correspond to a partition identifier. The stateful application can identify the new backup partition based on the hash value (operation **458**).

(102) FIG. 5A presents a flowchart illustrating the process of a partition management system downscaling a topic in an event-driven platform, in accordance with an aspect of the present application. During operation, the system can suspend the stateful application associated with a topic and a stateful backup topic defined on the event-driven platform such that the topic maintains a record log of a stream of computer-detectable events from the stateful application, and the stateful backup topic maintains the state of event records from the stateful application (operation **502**). The suspension can stop further updates from the stateful application to the event logs. Therefore, the system can perform the downscaling of the topic while the stateful application is suspended.

(103) To accommodate the downscaling of the topic to a reduced number of partitions (e.g., with a reduced partition count), the system can generate a placeholder for transferring the record logs and states in the topic and the stateful backup topic, respectively. Accordingly, the system can generate a temporary topic and a temporary stateful backup topic (operation **504**). These temporary topics can be generated with the reduced number of partitions. The system can regenerate the topic with a new set of partitions and the stateful backup topic with a new set of backup partitions at the event-driven platform (e.g., with the reduced number of partitions) (operation **506**).

(104) The system can determine a partition identifier indicating a new backup partition and a corresponding new backup partition for the stateful application (operation **508**). The system can use a hash function for the stateful application to determine the partition identifier. The system can then identify the new partition and the new backup partition allocated to the stateful application based on the partition identifier (operation **510**). The partition identifier can correspond to an index generated within a range indicated by the reduced number of partitions. The system can transfer the record log to the new partition via the temporary topic and the state to a new backup partition via the temporary stateful backup topic (operation **512**). Here, the temporary topic and the temporary stateful backup topic can operate as placeholders for the transfer.

(105) The stateful application can send subsequent event records to the event-driven platform. The new partition can allow the event-driven platform to store the subsequent event records from the stateful application in the new partition (operation **514**). The subsequent event records can be incorporated into an updated event log of the stateful application maintained in the new partition.

The system can resume the stateful application for generating an updated state, which is backed up at the new backup partition, based on the state and the subsequent event records associated with the stateful application (operation **516**).

(106) FIG. **5B** presents a flowchart illustrating the process of a stateful application distributing record logs and states to regenerated topics in an event-driven platform, in accordance with an aspect of the present application. During operation, the application can apply a hash function on a key associated with the stateful application to generate a hash value (operation **552**). The key can be an identifier of the stateful application. The hash value can be an index based on which a partition can be selected. Therefore, the index can correspond to a partition identifier. The stateful application can identify a temporary partition of the temporary topic and a temporary backup partition of the temporary stateful backup topic allocated to the stateful application based on the hash value (operation **554**).

(107) The system can then transfer the record log to the temporary partition and the state to the temporary backup partition (operation **556**). In this way, the system can preserve a copy of the record log and the state during the downscaling of the topic. Since the record log and the state are preserved, the system can remove the topic and the stateful backup topic from the event-driven platform upon transferring the event log and the state (operation **558**). The deletion can free the topic identifier at the event-driven platform. Therefore, the topic can be regenerated using the topic identifier.

(108) The system can generate a new topic at the event-driven platform such that the removed topic and the new topic have the same topic identifier (operation **560**). In the same way, the system can generate a new stateful backup topic at the event-driven platform such that the new stateful backup topic corresponds to the new topic (operation **562**). Because the new topic and the new stateful backup topic can be associated with the topic identifier, the stateful application does not need to be reconfigured. Accordingly, the system can transfer the event log to the new partition from the temporary partition upon regeneration of the topic (operation **564**). Similarly, the system can transfer the state to the new backup partition from the temporary backup partition upon regeneration of the stateful backup topic (operation **566**).

(109) FIG. **5C** presents a flowchart illustrating the process of a partition management system resuming consumer applications of a scaled topic in an event-driven platform, in accordance with an aspect of the present application. Here, the scaled topic can be either upscaled or downscaled. During operation, the system can suspend a consumer application associated with the topic (operation **572**). The suspension can stop the consumer application to retrieve data from the topic. Therefore, the system can perform the scaling of the topic while the consumer application is suspended.

(110) The system can check whether the state or event log transfer is complete (operation **574**). Until the transfer is complete, the system can continue the checking. If the transfer is complete, the system can resume the consumer application (operation **576**). The consumer application can then retrieve subsequent event records generated by the consumer application. The event records are incorporated into an event log in one of the new set of partitions of the topic. The new set of partitions can be the resultant partitions generated by the upscaling or downscaling of the topic. Accordingly, the system can allow the consumer application to receive the subsequent event records from one of the new set of partitions (operation **578**).

(111) FIG. **6** illustrates an example of a computing system facilitating dynamic topic repartitioning in an event-driven platform, in accordance with an aspect of the present application. A computing system **600** can include a set of processors **602**, a memory unit **604**, a network interface controller (NIC) **606**, and a storage device **608**. Memory unit **604** can include a set of volatile memory devices (e.g., dual in-line memory module (DIMM)). Furthermore, computing system **600** may be coupled to a display device **612**, a keyboard **614**, and a pointing device **616**, if needed. Computing system **600** can execute a partition management system **620**, which can operate in conjunction with

an event-driven platform, such as Kafka. Storage device **608** can store data **634** associated with partition management system **620**.

(112) Partition management system **620** can include instructions, which when executed by computing system **600**, can cause computing system **600** to perform methods and/or processes described in this disclosure. Partition management system **620** can operate in conjunction with an event management system. Furthermore, data **634** can include any data that can facilitate the operations of partition management system **620**. data **634** can include, but is not limited to, event records of a respective stateful application, record logs incorporating respective stateful of event records, states generated based on the corresponding stateful of event records, and backup copies of the states maintained by the event-driven platform.

(113) Specifically, partition management system **620** can include instructions for operating a proxy that can operate as a proxy between a user and partition management system **620** (e.g., via user interface) (proxy instructions **622**). Partition management system **620** can also include instructions for invoking a respective component of partition management system **620** (proxy instructions **622**). Partition management system **620** can further include instructions for suspending and resuming a respective stateful application associated with a topic defined on the event-driven platform (application instructions **624**).

(114) Moreover, partition management system **620** can include instructions for configuring the event-driven platform to upscale a topic (topic instructions **626**). Partition management system **620** can also include instructions for removing and generating topics and stateful backup topics at the event-driven platform (topic instructions **626**). Furthermore, partition management system **620** can include instructions for publishing or transferring event logs to a partition of a topic (publishing instructions **628**). Partition management system **620** can also include instructions for publishing or transferring states to a backup partition of a stateful backup topic (publishing instructions **628**).

(115) In addition, partition management system **620** can include instructions for generating a partition identifier that can identify a partition or backup partition associated with a stateful application, a record log, or a state (partition instructions **630**). Partition management system **620** can also include instructions for applying a hash function on a key associated with a stateful application, a record log, or a state for generating an index, which can correspond to the partition identifier (partition instructions **630**). Furthermore, partition management system **620** can include instructions for sending and receiving messages based on respective topics (communication instructions **632**).

(116) FIG. 7 illustrates an example of a CRM that facilitates dynamic topic repartitioning in an event-driven platform, in accordance with an aspect of the present application. CRM **700** can be a non-transitory computer-readable medium or device storing instructions that when executed by a computer or processor cause the computer or processor to perform a method. CRM **700** can include proxy instructions **702** for operating a proxy that can operate as a proxy between a user and a partition management system (e.g., via user interface) and invoking a respective component of the partition management system. CRM **700** can include application instructions **704** for suspending and resuming a respective stateful application associated with a topic defined on an event-driven platform.

(117) CRM **700** can also include topic instructions **706** for configuring the event-driven platform to upscale a topic and removing and generating topics and stateful backup topics at the event-driven platform. Furthermore, CRM **700** can include publishing instructions **708** for publishing or transferring event logs to a partition of a topic and states to a backup partition of a stateful backup topic. Moreover, CRM **700** can include partition instructions **710** for generating a partition identifier that can identify a partition or backup partition associated with a stateful application, a record log, or a state. Partition instructions **710** can also be for applying a hash function on a key associated with a stateful application, a record log, or a state for generating an index, which can correspond to the partition identifier. CRM **700** can also include communication instructions **712**

for sending and receiving messages based on respective topics.

(118) The description herein is presented to enable any person skilled in the art to make and use the invention and is provided in the context of a particular application and its requirements. Various modifications to the disclosed examples will be readily apparent to those skilled in the art, and the general principles defined herein may be applied to other examples and applications without departing from the spirit and scope of the present invention. Thus, the present invention is not limited to the examples shown but is to be accorded the widest scope consistent with the claims.

(119) One aspect of the present technology can provide a system facilitating scaling of topics on an event-driven platform. During operation, the system can suspend a stateful application associated with a topic and a stateful backup topic defined on the event-driven platform. The topic can correspond to a stream of computer-detectable events associated with the stateful application operating in conjunction with the event-driven platform. The system can then modify the configuration of the topic at the event-driven platform to allow a new set of partitions associated with the topic. The modified configuration generates a new set of backup partitions for the stateful backup topic. The system can identify a state generated based on event records from the stateful application in a previous backup partition of the stateful backup topic and identify a new backup partition allocated to the stateful application. If the previous and new backup partitions are different, the system can transfer the state from the previous backup partition to the new backup partition. The system can then resume the stateful application for generating an updated state, which is maintained at the new backup partition, based on the state and subsequent event records associated with the stateful application.

(120) In an example embodiment, the system can identify the new backup partition by determining a partition identifier associated with the state. The partition identifier can correspond to the new backup partition. The system can determine whether the previous and new backup partitions are different based on corresponding partition identifiers.

(121) In an example embodiment, the system can allow the event-driven platform to retrieve the state from the new backup partition to a new partition corresponding to the new backup partition.

(122) In an example embodiment, the system can allow the event-driven platform to store the subsequent event records from the stateful application in the new partition.

(123) In an example embodiment, the system can allow the event-driven platform to transfer the event records from the previous partition to the new partition in response to receiving a subsequent event record from the stateful application.

(124) In an example embodiment, the system can identify the new backup partition by applying a hash function on a key associated with the stateful application to generate a hash value and identifying the new backup partition based on the hash value.

(125) In an example embodiment, the hash function can be used to determine a new partition for storing event records from the stateful application.

(126) In an example embodiment, the previous and new backup partitions can be generated based on a same partitioning method.

(127) In an example embodiment, the system can suspend a consumer application associated with the topic. Upon transferring the state from the previous backup partition to the new backup partition, the system can resume the consumer application for receiving subsequent event records from one of the new set of partitions.

(128) In an example embodiment, a number of new backup partitions is greater than a number of previous backup partitions of the stateful backup topic.

(129) Another aspect of the present technology can provide a system facilitating scaling of topics on an event-driven platform. During operation, the system can suspend a stateful application associated with a topic and a stateful backup topic defined on an event-driven platform. A partition of the topic maintains an event log of a stream of computer-detectable events associated with the stateful application operating in conjunction with the event-driven platform. The stateful backup

topic maintains a state generated based on event records from the stateful application. The system can generate a temporary topic and a temporary stateful backup topic the event-driven platform. The system can then regenerate, at the event-driven platform, the topic with a new set of partitions and the stateful backup topic with a new set of backup partitions. The system can also identify a new partition and a new backup partition allocated to the stateful application and transfer the event log to the new partition via the temporary topic and the state to the new backup partition via the temporary stateful backup topic. The system can then resume the stateful application for generating an updated state, which is maintained at the new backup partition, based on the state and subsequent event records associated with the stateful application.

(130) In an example embodiment, the system can identify the new partition by determining a partition identifier associated with the stateful application. Here, the partition identifier can correspond to the new partition.

(131) In an example embodiment, the system can identify a temporary partition of the temporary topic and a temporary backup partition of the temporary stateful backup topic allocated to the stateful application. The system can then transfer the event log to the temporary partition and the state to the temporary backup partition.

(132) In an example embodiment, the system can regenerate the topic by, upon transferring the event log and the state, removing the topic and the stateful backup topic from the event-driven platform. The system can generate, at the event-driven platform, a new topic such that the removed topic and the new topic have a same topic identifier. The system can also generate, at the event-driven platform, a new stateful backup topic such that the new stateful backup topic corresponds to the new topic.

(133) In an example embodiment, upon regenerating the topic, the system can transfer the event log to the new partition from the temporary partition. Similarly, upon regenerating the stateful backup topic, the system can transfer the state to the new backup partition from the temporary backup partition.

(134) In an example embodiment, the system can identify the temporary partition by applying a hash function on a key associated with the stateful application to generate a hash value and identify the temporary partition based on the hash value.

(135) In an example embodiment, the temporary and new partitions are generated based on a same partitioning method.

(136) In an example embodiment, the system can allow the event-driven platform to store the subsequent event records from the stateful application in the new partition.

(137) In an example embodiment, the system can suspend a consumer application associated with the topic. Upon transferring the event log to the new partition, the system can resume the consumer application for receiving subsequent event records from one of the new set of partitions.

(138) In an example embodiment, a number of new partitions in the regenerated topic is less than a number of previous partitions of the topic.

(139) The data structures and code described in this detailed description are typically stored on a computer-readable storage medium, which may be any device or medium that can store code and/or data for use by a computer system. The computer-readable storage medium includes, but is not limited to, volatile memory, non-volatile memory, magnetic and optical storage devices such as disks, magnetic tape, CDs (compact discs), DVDs (digital versatile discs or digital video discs), or other media capable of storing computer-readable media now known or later developed.

(140) The methods and processes described in the detailed description section can be embodied as code and/or data, which can be stored in a computer-readable storage medium as described above. When a computer system reads and executes the code and/or data stored on the computer-readable storage medium, the computer system performs the methods and processes embodied as data structures and code and stored within the computer-readable storage medium.

(141) The methods and processes described herein can be executed by and/or included in hardware

logic blocks or apparatus. These logic blocks or apparatus may include, but are not limited to, an application-specific integrated circuit (ASIC) chip, a field-programmable gate array (FPGA), a dedicated or shared processor that executes a particular software logic block or a piece of code at a particular time, and/or other programmable-logic devices now known or later developed. When the hardware logic blocks or apparatus are activated, they perform the methods and processes included within them.

(142) The foregoing descriptions of examples of the present invention have been presented only for purposes of illustration and description. They are not intended to be exhaustive or to limit this disclosure. Accordingly, many modifications and variations will be apparent to practitioners skilled in the art. The scope of the present invention is defined by the appended claims.

Claims

1. A computer-implemented method, comprising: suspending a stateful application associated with a topic and a stateful backup topic defined on an event-driven platform, wherein the topic corresponds to a stream of computer-detectable events associated with the stateful application operating in conjunction with the event-driven platform; modifying configuration of the topic at the event-driven platform to allow a new set of partitions associated with the topic, wherein the modified configuration generates a new set of backup partitions for the stateful backup topic; identifying a state generated based on event records from the stateful application in a previous backup partition of the stateful backup topic; identifying a new backup partition allocated to the stateful application; in response to the previous and new backup partitions being different, transferring the state from the previous backup partition to the new backup partition; and resuming the stateful application for generating an updated state, which is maintained at the new backup partition, based on the state and subsequent event records associated with the stateful application.
2. The method of claim 1, wherein identifying the new backup partition further comprises determining a partition identifier associated with the state, wherein the partition identifier corresponds to the new backup partition; and wherein the method further comprises determining whether the previous and new backup partitions are different based on corresponding partition identifiers.
3. The method of claim 1, further comprising allowing the event-driven platform to retrieve the state from the new backup partition to a new partition corresponding to the new backup partition.
4. The method of claim 3, further comprising allowing the event-driven platform to store the subsequent event records from the stateful application in the new partition.
5. The method of claim 3, further comprising allowing the event-driven platform to transfer the event records from the previous partition to the new partition in response to receiving a subsequent event record from the stateful application.
6. The method of claim 1, wherein identifying the new backup partition comprises: applying a hash function on a key associated with the stateful application to generate a hash value; and identifying the new backup partition based on the hash value.
7. The method of claim 6, wherein the hash function is used to determine a new partition for storing event records from the stateful application.
8. The method of claim 1, wherein the previous and new backup partitions are generated based on a same partitioning method.
9. The method of claim 1, further comprising: suspending a consumer application associated with the topic; and in response to transferring the state from the previous backup partition to the new backup partition, resuming the consumer application for receiving subsequent event records from one of the new set of partitions.
10. The method of claim 1, wherein a number of new backup partitions is greater than a number of previous backup partitions of the stateful backup topic.

11. A computer-implemented method, comprising: suspending a stateful application associated with a topic and a stateful backup topic defined on an event-driven platform, wherein a partition of the topic maintains an event log of a stream of computer-detectable events associated with the stateful application operating in conjunction with the event-driven platform, and wherein the stateful backup topic maintains a state generated based on event records from the stateful application; generating a temporary topic and a temporary stateful backup topic at the event-driven platform; regenerating, at the event-driven platform, the topic with a new set of partitions and the stateful backup topic with a new set of backup partitions; identifying a new partition and a new backup partition allocated to the stateful application; transferring the event log to the new partition via the temporary topic and the state to the new backup partition via the temporary stateful backup topic; and resuming the stateful application for generating an updated state, which is maintained at the new backup partition, based on the state and subsequent event records associated with the stateful application.

12. The method of claim 11, wherein identifying the new partition further comprises determining a partition identifier associated with the stateful application, wherein the partition identifier corresponds to the new partition.

13. The method of claim 11, further comprising: identifying a temporary partition of the temporary topic and a temporary backup partition of the temporary stateful backup topic allocated to the stateful application; and transferring the event log to the temporary partition and the state to the temporary backup partition.

14. The method of claim 13, wherein regenerating the topic further comprises: in response to transferring the event log and the state, removing the topic and the stateful backup topic from the event-driven platform; generating, at the event-driven platform, a new topic such that the removed topic and the new topic have a same topic identifier; and generating, at the event-driven platform, a new stateful backup topic such that the new stateful backup topic corresponds to the new topic.

15. The method of claim 14, further comprising: in response to regenerating the topic, transferring the event log to the new partition from the temporary partition; and in response to regenerating the stateful backup topic, transferring the state to the new backup partition from the temporary backup partition.

16. The method of claim 13, wherein identifying the temporary partition comprises: applying a hash function on a key associated with the stateful application to generate a hash value; and identifying the temporary partition based on the hash value.

17. The method of claim 13, wherein the temporary and new partitions are generated based on a same partitioning method.

18. The method of claim 11, further comprising allowing the event-driven platform to store the subsequent event records from the stateful application in the new partition.

19. The method of claim 11, further comprising: suspending a consumer application associated with the topic; and in response to transferring the event log to the new partition, resuming the consumer application for receiving subsequent event records from one of the new set of partitions.

20. A computer system, comprising: a processing resource; and a non-transitory computer-readable storage medium storing instructions that when executed by the processing resource cause the computer system to: suspend a stateful application associated with a topic and a stateful backup topic defined on an event-driven platform, wherein the topic corresponds to a stream of computer-detectable events associated with the stateful application operating in conjunction with the event-driven platform; modify configuration of the topic at the event-driven platform to allow a new set of partitions associated with the topic, wherein the modified configuration generates a new set of backup partitions for the stateful backup topic; identify a state generated based on event records from the stateful application in a previous backup partition of the stateful backup topic; identify a new backup partition allocated to the stateful application; in response to the previous and new backup partitions being different, transfer the state from the previous backup partition to the new

backup partition; and resume the stateful application for generating an updated state, which is maintained at the new backup partition, based on the state and subsequent event records associated with the stateful application.
