

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12405866
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Sity; Elad et al.

High performance processor for low-way and high-latency memory instances

Abstract

Distributed processors and methods for compiling code for execution by distributed processors are disclosed. In one implementation, a distributed processor may include a substrate; a memory array disposed on the substrate; and a processing array disposed on the substrate. The memory array may include a plurality of discrete memory banks, and the processing array may include a plurality of processor subunits, each one of the processor subunits being associated with a corresponding, dedicated one of the plurality of discrete memory banks. The distributed processor may further include a first plurality of buses, each connecting one of the plurality of processor subunits to its corresponding, dedicated memory bank, and a second plurality of buses, each connecting one of the plurality of processor subunits to another of the plurality of processor subunits.

Inventors:	Sity; Elad (Kfar Saba, IL), Hillel; Eliad (Kfar Saba, IL)
Applicant:	NeuroBlade Ltd. (Tel Aviv, IL)
Family ID:	65233543
Assignee:	NeuroBlade Ltd. (Tel Aviv, IL)
Appl. No.:	18/406346
Filed:	January 08, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240143457 A1	May. 02, 2024

Related U.S. Application Data

continuation parent-doc US 17397061 20210809 US 11914487 child-doc US 18406346
continuation parent-doc WO PCT/IB2018/000995 20180730 PENDING child-doc US 16512622

division parent-doc US 16512622 20190716 US 11126511 20210921 child-doc US 17397061
us-provisional-application US 62548990 20170823
us-provisional-application US 62538722 20170730
us-provisional-application US 62538724 20170730

Publication Classification

Int. Cl.: G06F9/38 (20180101); G06F9/30 (20180101); G06F11/10 (20060101); G06F11/16 (20060101); G06F13/16 (20060101); G06F15/80 (20060101); G06N3/04 (20230101); G11C7/10 (20060101); G11C11/16 (20060101); G11C11/4076 (20060101); G11C11/408 (20060101); G11C11/4093 (20060101); G06F15/76 (20060101)

U.S. Cl.:

CPC G06F11/1658 (20130101); G06F9/3001 (20130101); G06F9/3885 (20130101); G06F9/3889 (20130101); G06F9/3895 (20130101); G06F11/1016 (20130101); G06F11/102 (20130101); G06F11/16 (20130101); G06F11/1616 (20130101); G06F13/1657 (20130101); G06F15/8038 (20130101); G06N3/04 (20130101); G11C7/1072 (20130101); G11C11/1655 (20130101); G11C11/1657 (20130101); G11C11/1675 (20130101); G11C11/4076 (20130101); G11C11/408 (20130101); G11C11/4093 (20130101); G06F2015/765 (20130101)

Field of Classification Search

CPC: G06F (13/1657)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5155729	12/1991	Rysko et al.	N/A	N/A
5179702	12/1992	Spix et al.	N/A	N/A
5214747	12/1992	Cok	N/A	N/A
5239654	12/1992	Ing-Simmons et al.	N/A	N/A
5297260	12/1993	Kametani	N/A	N/A
5345552	12/1993	Brown	N/A	N/A
5502728	12/1995	Smith, III	N/A	N/A
5506992	12/1995	Saxenmeyer	N/A	N/A
5517600	12/1995	Shimokawa	N/A	N/A
5568617	12/1995	Kametani	N/A	N/A
5590345	12/1995	Barker et al.	N/A	N/A
5625796	12/1996	Kaczmarczyk et al.	N/A	N/A
5710935	12/1997	Barker et al.	N/A	N/A
5751987	12/1997	Mahant-Shetti et al.	N/A	N/A
5752036	12/1997	Nakamura et al.	N/A	N/A
5752067	12/1997	Wikinson et al.	N/A	N/A
5822608	12/1997	Dieffenderfer et al.	N/A	N/A
5956703	12/1998	Turner et al.	N/A	N/A
6026464	12/1999	Cohen	N/A	N/A

6041400	12/1999	Ozcelik et al.	N/A	N/A
6064621	12/1999	Tanizaki et al.	N/A	N/A
6067260	12/1999	Ooishi et al.	N/A	N/A
6067262	12/1999	Irrinki et al.	N/A	N/A
6096094	12/1999	Kay et al.	N/A	N/A
6145069	12/1999	Dye	N/A	N/A
6216178	12/2000	Stracovsky et al.	N/A	N/A
6363453	12/2001	Esposito et al.	N/A	N/A
6389497	12/2001	Koslowsky et al.	N/A	N/A
6396760	12/2001	Behera et al.	N/A	N/A
6449732	12/2001	Rasmussen et al.	N/A	N/A
6453398	12/2001	McKenzie	N/A	N/A
6523018	12/2002	Louis et al.	N/A	N/A
6553355	12/2002	Arnoux et al.	N/A	N/A
6584033	12/2002	Ayukawa et al.	N/A	N/A
6601126	12/2002	Zaidi et al.	N/A	N/A
6640283	12/2002	Naffziger et al.	N/A	N/A
6668308	12/2002	Barroso et al.	N/A	N/A
6678801	12/2003	Greim et al.	N/A	N/A
6708254	12/2003	Lee et al.	N/A	N/A
6717834	12/2003	Zagorianakos et al.	N/A	N/A
6785780	12/2003	Klein et al.	N/A	N/A
6785841	12/2003	Akrout et al.	N/A	N/A
6798420	12/2003	Xie	348/715	G09G 5/346
6823429	12/2003	Olnowich	N/A	N/A
6877046	12/2004	Singh	N/A	N/A
6894392	12/2004	Gudesen et al.	N/A	N/A
6988154	12/2005	Latta	N/A	N/A
6999950	12/2005	Linneberg et al.	N/A	N/A
7107285	12/2005	von Kaenel et al.	N/A	N/A
7107412	12/2005	Klein et al.	N/A	N/A
7159141	12/2006	Lakhani et al.	N/A	N/A
7194568	12/2006	Jeter, Jr. et al.	N/A	N/A
7290109	12/2006	Horii et al.	N/A	N/A
7325232	12/2007	Liem	N/A	N/A
7450410	12/2007	Klein	N/A	N/A
7549081	12/2008	Robbins et al.	N/A	N/A
7557605	12/2008	D'Souza et al.	N/A	N/A
7558934	12/2008	Kondo et al.	N/A	N/A
7562271	12/2008	Shaeffer et al.	N/A	N/A
7610537	12/2008	Dickinson et al.	N/A	N/A
7657712	12/2009	Lentz et al.	N/A	N/A
7772880	12/2009	Solomon	N/A	N/A
7782703	12/2009	Oh	N/A	N/A
7836168	12/2009	Vasko et al.	N/A	N/A
7881321	12/2010	Deneroff et al.	N/A	N/A
7882307	12/2010	Wentzlaff et al.	N/A	N/A
7882320	12/2010	Caulkins	N/A	N/A
7949820	12/2010	Caulkins	N/A	N/A

8028124	12/2010	Ban et al.	N/A	N/A
8042082	12/2010	Solomon	N/A	N/A
8055601	12/2010	Pandya	N/A	N/A
8074031	12/2010	Bekooij	N/A	N/A
8171233	12/2011	Kwon	N/A	N/A
8369123	12/2012	Kang et al.	N/A	N/A
8442927	12/2012	Chakradhar et al.	N/A	N/A
8648403	12/2013	Luk et al.	N/A	N/A
8677306	12/2013	Andreev et al.	N/A	N/A
8738860	12/2013	Griffin et al.	N/A	N/A
8984256	12/2014	Fish	N/A	N/A
9003109	12/2014	Lam	N/A	N/A
9053762	12/2014	Hirobe	N/A	N/A
9098209	12/2014	Gopalakrishnan et al.	N/A	N/A
9164807	12/2014	Blanc et al.	N/A	N/A
9177611	12/2014	D'Abreu	N/A	G06F 11/1048
9239691	12/2015	Lam	N/A	N/A
9245222	12/2015	Modha	N/A	N/A
9305635	12/2015	Lin et al.	N/A	N/A
9342474	12/2015	Macri et al.	N/A	N/A
9348385	12/2015	De Rochemont et al.	N/A	N/A
9378003	12/2015	Sundararajan et al.	N/A	N/A
9432298	12/2015	Smith	N/A	N/A
9449257	12/2015	Shi et al.	N/A	N/A
9477636	12/2015	Walker et al.	N/A	N/A
9570132	12/2016	Kim et al.	N/A	N/A
9653151	12/2016	Ong et al.	N/A	N/A
9665503	12/2016	Dalal	N/A	N/A
9672169	12/2016	Winderweedle	N/A	N/A
9760827	12/2016	Lin et al.	N/A	N/A
9904874	12/2017	Shoaib et al.	N/A	N/A
9978014	12/2017	Lupon et al.	N/A	N/A
9984337	12/2017	Kadav et al.	N/A	N/A
10032110	12/2017	Young et al.	N/A	N/A
10067681	12/2017	Park et al.	N/A	N/A
10078620	12/2017	Farabet et al.	N/A	N/A
10114795	12/2017	Cargnini et al.	N/A	N/A
10140252	12/2017	Fowers et al.	N/A	N/A
10157309	12/2017	Molchanov et al.	N/A	N/A
10423876	12/2018	Henry	N/A	G06F 12/0811
2005/0138325	12/2004	Hofstee et al.	N/A	N/A
2005/0223275	12/2004	Jardine et al.	N/A	N/A
2007/0022063	12/2006	Lightowler	N/A	N/A
2007/0214335	12/2006	Bellows et al.	N/A	N/A
2007/0220369	12/2006	Fayad et al.	N/A	N/A
2007/0276977	12/2006	Coteus et al.	N/A	N/A
2008/0046665	12/2007	Kim	N/A	N/A
2009/0083263	12/2008	Felch et al.	N/A	N/A

2009/0158010	12/2008	Gartner et al.	N/A	N/A
2009/0196116	12/2008	Oh	N/A	N/A
2009/0217222	12/2008	Yasuda et al.	N/A	N/A
2010/0131827	12/2009	Sokolov et al.	N/A	N/A
2011/0016278	12/2010	Ware et al.	N/A	N/A
2011/0296078	12/2010	Khan et al.	N/A	N/A
2011/0314473	12/2010	Yang et al.	N/A	N/A
2012/0102576	12/2011	Chew	N/A	N/A
2012/0255018	12/2011	Sallam	N/A	N/A
2012/0291131	12/2011	Nagpal et al.	N/A	N/A
2013/0339589	12/2012	Qawami	711/103	G11C 13/0004
2014/0040622	12/2013	Kendall et al.	N/A	N/A
2014/0136915	12/2013	Hyde et al.	N/A	N/A
2014/0143470	12/2013	Dobbs et al.	N/A	N/A
2014/0310495	12/2013	Michelogiannakis et al.	N/A	N/A
2014/0328103	12/2013	Arsovski	N/A	N/A
2014/0359199	12/2013	Huppenthal et al.	N/A	N/A
2015/0089162	12/2014	Ahsan et al.	N/A	N/A
2015/0146491	12/2014	Akerib	365/189.02	G11C 7/18
2015/0212861	12/2014	Canoy et al.	N/A	N/A
2015/0293859	12/2014	Huang et al.	N/A	N/A
2015/0309779	12/2014	Baskaran et al.	N/A	N/A
2015/0324690	12/2014	Chilimbi et al.	N/A	N/A
2016/0034809	12/2015	Trenholm et al.	N/A	N/A
2016/0232445	12/2015	Srinivasan et al.	N/A	N/A
2016/0260024	12/2015	Campos et al.	N/A	N/A
2016/0379109	12/2015	Chung et al.	N/A	N/A
2017/0103298	12/2016	Ling et al.	N/A	N/A
2017/0194045	12/2016	Kang	N/A	N/A
2017/0200094	12/2016	Bruestle et al.	N/A	N/A
2017/0286825	12/2016	Akopyan et al.	N/A	N/A
2017/0301389	12/2016	Kajigaya	N/A	G06F 13/161
2018/0052766	12/2017	Mehra et al.	N/A	N/A
2018/0144244	12/2017	Masoud et al.	N/A	N/A
2018/0189215	12/2017	Boesch et al.	N/A	N/A
2018/0189645	12/2017	Chen et al.	N/A	N/A
2018/0210830	12/2017	Malladi et al.	N/A	N/A
2018/0285254	12/2017	Baum et al.	N/A	N/A
2018/0285718	12/2017	Baum et al.	N/A	N/A
2018/0336035	12/2017	Choi et al.	N/A	N/A
2019/0102325	12/2018	Natu et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
2 149 479	12/2000	CA	N/A

OTHER PUBLICATIONS

International Searching Authority, Notification of Transmittal of the International Search Report and the Written Opinion of the International Search Authority, International Application No. PCT/IB2020/00665, dated May 27, 2021, 15 pages. cited by applicant

Ahn et al., "A Scalable Processing-in-Memory Accelerator for Parallel Graph Processing," ISCA '15 (Jun. 13-17, 2015), pp. 105-117. cited by applicant

Stripf et al., "Compiling Scilab to high performance embedded multicore systems," Microprocessors and Microsystems (2013), 37:1033-49. cited by applicant

Chandramoorthy, "Design and Exploration of Accelerator-Rich Multi Core Architectures," The Pennsylvania State University (Dec. 2016), pp. title page, ii-xii, and 1-106. cited by applicant

Chaturvedi, "Techniques to Improve the Performance of Cache Memory for Multi-Core Processors," Birla Institute of Technology and Science (2015), pp. title page, i-xv, and 1-185. cited by applicant

Chetlur et al., "cuDNN: Efficient Primitives for Deep Learning," arXiv:1410.0759v3 (Dec. 18, 2014), pp. 1-9. cited by applicant

Chi et al., "Prime: A Novel Processing-in-memory Architecture for Neural Network Computation in ReRAM-based Main Memory," 2016 ACM/IEEE 43rd Annual International Symposium on Computer Architecture (2016), pp. 27-39. cited by applicant

Chung, "CoRAM: An In-Fabric Memory Architecture for FPGA-Based Computing," Carnegie Mellon University (Aug. 2011), pp. title page, i-xiv, 1-179. cited by applicant

Cristian, "Understanding Fault-Tolerant Distributed Systems," University of California, San Diego (May 25, 1993), pp. 1-45. cited by applicant

De Dinechin et al., "A Clustered Manycore Processor Architecture for Embedded and Accelerated Applications," ResearchGate (Jan. 2013), pp. 1-6. cited by applicant

Gupta et al., "The StageNet fabric for constructing resilient multicore systems," Proceedings of the 41.SUP.st .annual IEEE/ACM International Symposium on Microarchitecture (Nov. 8-12, 2008), Abstract page. cited by applicant

HP Proliant Servers Troubleshooting Guide, Hewlett Packard (Oct. 30, 2011), Edition 12. cited by applicant

Draper et al., "Implementation of a 32-bit RISC Processor for the Data-Intensive Architecture Processing-In-Memory Chip," Proceedings of the IEEE International Conference on Application-Specific Systems, Architectures, and Processors (2002), pp. 1-10. cited by applicant

Malki et al., "A CNN-Specific Integrated Processor," EURASIP Journal on Advances in Signal Processing (2009), 2009: pp. 1-14. cited by applicant

Meridian 1 Engineering Handbook, Update Package, Meridian 1 Communication Systems (Apr. 30, 1990), pp. 1-642. cited by applicant

Nanyang Technological University, "Scientists turn memory chips into processors to speed up computing tasks," Nanyang Technological University (Jan. 3, 2017), pp. 1-5. cited by applicant

Shafiee et al., "ISAAC: A Convolutional Neural Network Accelerator with In-Situ Analog Arithmetic in Crossbars," ACM SIGARCH Computer Architecture News 44 (2016), pp. 14-26. cited by applicant

Weninger et al., "Introducing CURRENNT: The Munich Open-Source CUDA RecurREnt Neural Network Toolkit," Journal of Machine Learning Research (Mar. 2015), pp. 547-551. cited by applicant

Wu et al., "Demonstration and architectural analysis of complementary metal-oxide semiconductor/multiple-quantum-well smart-pixel array cellular logic processors for single-instruction multiple-data parallel-pipeline processing," Applied Optics (Apr. 10, 1999), 38:2270-81. cited by applicant

Yuan et al., "Complexity Effective Memory Access Scheduling for Many-Core Accelerator Architectures," MICRO '09 (Dec. 12-16, 2009), pp. 34-44. cited by applicant

Udipi et al., “Rethinking DRAM Design and Organization for Energy-Constrained Multi-Cores,” ISCA '10 (Jun. 19-23, 2010), pp. 175-186. cited by applicant
Chen et al., “Eyeriss: A Spatial Architecture for Energy-Efficient Dataflow for Convolutional Neural Networks,” 2016 ACM/IEEE 43.SUP.rd .Annual International Symposium on Computer Architecture (2016), pp. 1-13. cited by applicant
Gao et al., “TETRIS: Scalable and Efficient Neural Network Acceleration with 3D Memory,” ASPLOS '17 (Apr. 8-12, 2017), pp. 1-14. cited by applicant
Chen et al., “DianNao: A Small-Footprint High-Throughput Accelerator for Ubiquitous Machine-Learning,” ASPLOS '14 (Mar. 1-5, 2014), pp. 1-1-15. cited by applicant
Azarkhish et al., “Neurostream: Scalable and Energy Efficient Deep Learning with Smart Memory Cubes,” arXiv:1701.06420v3 (Sep. 24, 2017), pp. 1-15. cited by applicant
Farabet et al., “NeuFlow: A Runtime Reconfigurable Dataflow Processor Vision,” CVPR 2011 Workshops (Jun. 20-25, 2011), pp. 109-116. cited by applicant
Wang et al., “DLAU: A Scalable Deep Learning Accelerator Unit on Fpga,” IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems (2016), pp. 1-5. cited by applicant
Han et al., “EIE: Efficient Inference Engine on Compressed Deep Neural Network,” arXiv:1602.01528v2 (May 3, 2016), pp. 1-12. cited by applicant
European Search Report mailed Apr. 28, 2023, by the European Patent Office in European Application No. EP 23 15 1586, 15 pages. cited by applicant

Primary Examiner: Metzger; Michael J

Attorney, Agent or Firm: Finnegan, Henderson, Farabow, Garrett & Dunner LLP

Background/Summary

CROSS REFERENCES TO RELATED APPLICATIONS (1) This application is a continuation of U.S. application Ser. No. 17/397,061, filed on Aug. 9, 2021 (now allowed), which is a divisional of U.S. application Ser. No. 16/512,622, filed on Jul. 16, 2019, issued as U.S. Pat. No. 11,126,511 on Sep. 21, 2021, which is a continuation of PCT/IB2018/000995, filed on Jul. 30, 2018, which claims the benefit of priority of U.S. Provisional Patent Application No. 62/538,722, filed on Jul. 30, 2017, U.S. Provisional Patent Application No. 62/538,724, filed on Jul. 30, 2017, and U.S. Provisional Patent Application No. 62/548,990, filed on Aug. 23, 2017. All of the foregoing applications are incorporated herein by reference in their entireties.

BACKGROUND

Technical Field

(1) The present disclosure relates generally to apparatuses for facilitating memory-intensive operations. In particular, the present disclosure relates to hardware chips that include processing elements coupled to dedicated memory banks.

Background Information

(2) As processor speeds and memory sizes both continue to increase, a significant limitation on effective processing speeds is the von Neumann bottleneck. The von. Neumann bottleneck results from throughput limitations resulting from conventional computer architecture. In particular, data transfer from memory to the processor is often bottlenecked compared to actual computations undertaken by the processor. Accordingly, the number of clock cycles to read and write from memory increases significantly with memory-intensive processes. These clock cycles result in lower effective processing speeds because reading and writing from memory consumes clock cycles that cannot be used for performing operations on data. Moreover, the computational

bandwidth of the processor is generally larger than the bandwidth of the buses that the processor uses to access the memory.

(3) These bottlenecks are particularly pronounced for memory-intensive processes, such as neural network and other machine learning algorithms; database construction, indexing searching, and querying; and other tasks that include more reading and writing operation than data processing operations.

(4) Additionally, the rapid growth in volume and granularity of available digital data has created opportunities to develop machine learning algorithms and has enabled new technologies. However, it has also brought cumbersome challenges to the world of data bases and parallel computing. For example, the rise of social media and the Internet of Things (IoT) creates digital data at a record rate. This new data can be used to create algorithms for a variety of purposes, ranging from new advertising techniques to more precise control methods of industrial processes. However, the new data has been difficult to store, process, analyze and handle.

(5) New data resources can be massive, sometimes in the order of peta- to zettabytes. Moreover, the growth rate of these data resources may exceed data processing capabilities. Therefore, data scientists have turned to parallel data processing techniques, to tackle these challenges. In an effort to increase computation power and handle the massive amount of data, scientists have attempted to create systems and methods capable of parallel intensive computing. But these existing systems and methods have not kept up with the data processing requirements, often because the techniques employed are limited by their demand of additional resources for data management, integration of segregated data, and analysis of the sectioned data.

(6) To facilitate the manipulation of large data sets, engineers and scientists now seek to improve the hardware used to analyze data. For example, new semiconductor processors or chips (such as those described herein) may be designed specifically for data intensive tasks by incorporating memory and processing functions in a single substrate fabricated in technologies more fitting for memory operations rather than arithmetic computation. With integrated circuits specifically designed for data-intensive tasks, it is possible to meet the new data processing requirements. Nonetheless, this new approach to tackle data processing of large data sets requires solving new issues in chip design and fabrication. For instance, if the new chips designed for data intensive tasks are manufactured with fabrication techniques and architectures used for common chips, they would have a poor performance and/or unacceptable yields. In addition, if the new chips are designed to operate with current data handling methods, they will have poor performance because current methods can limit the chip's ability to handle parallel operations.

(7) The present disclosure describes solutions for mitigating or overcoming one or more of the problems set forth above, among other problems in the prior art.

SUMMARY

(8) Embodiments consistent with the present disclosure provide apparatuses including hardware processing chips. The disclosed embodiments may use dedicated memory banks paired with processing elements to provide more efficient effective processing speeds than conventional processors. For example, consistent with the disclosed embodiments, the disclosed chips may include dedicated buses between each processing element and its corresponding memory banks. In addition, the disclosed chips may be free of arbiters and/or other hardware that controls timing of data transfers between the processing elements. Other disclosed non-transitory computer-readable media may store instructions for compiling higher-level instructions to lower-level instructions executed by hardware chips disclosed herein.

(9) Some embodiments of the present disclosure include a distributed processor, comprising: a substrate; a memory array disposed on the substrate, the memory array including a plurality of discrete memory banks; a processing array disposed on the substrate, the processing array including a plurality of processor subunits, each one of the processor subunits being associated with a corresponding, dedicated one of the plurality of discrete memory banks; a first plurality of buses,

each connecting one of the plurality of processor subunits to its corresponding, dedicated memory bank; and a second plurality of buses, each connecting one of the plurality of processor subunits to another of the plurality of processor subunits.

(10) Other embodiments consistent with the present disclosure include a memory chip, comprising: a substrate; a memory array disposed on the substrate, the memory array including a plurality of discrete memory banks; a processing array disposed on the substrate, the processing array including a plurality of address generators, each one of the address generators being associated with a corresponding, dedicated one of the plurality of discrete memory banks; and a plurality of buses, each connecting one of the plurality of address generators to its corresponding, dedicated memory bank.

(11) Another embodiment consistent with the present disclosure may include a distributed processor, comprising: a substrate; a memory array disposed on the substrate, the memory array including a plurality of discrete memory banks, wherein each of the discrete memory banks has a capacity greater than one megabyte; and a processing array disposed on the substrate, the processing array including a plurality of processor subunits, each one of the processor subunits being associated with a corresponding, dedicated one of the plurality of discrete memory banks.

(12) Still other embodiments consistent with the present disclosure may include a distributed processor, comprising: a substrate; a memory array disposed on the substrate, the memory array including a plurality of discrete memory banks; and a processing array disposed on the substrate, the processing array including a plurality of processor subunits, each one of the processor subunits being associated with a corresponding, dedicated one of the plurality of discrete memory banks; and a plurality of buses, each one of the plurality of buses connecting one of the plurality of processor subunits to at least another one of the plurality of processor subunits, wherein the plurality of buses are free of timing hardware logic components such that data transfers between processor subunits and across corresponding ones of the plurality of buses are uncontrolled by timing hardware logic components.

(13) Other embodiments may include a distributed processor on a memory chip, comprising: a substrate; a memory array disposed on the substrate, the memory array including a plurality of discrete memory banks; and a processing array disposed on the substrate, the processing array including a plurality of processor subunits, each one of the processor subunits being associated with a corresponding, dedicated one of the plurality of discrete memory banks; and a plurality of buses, each one of the plurality of buses connecting one of the plurality of processor subunits to a corresponding, dedicated one of the plurality of discrete memory banks, wherein the plurality of buses are free of timing hardware logic components such that data transfers between a processor subunit and a corresponding, dedicated one of the plurality of discrete memory banks and across a corresponding one of the plurality of buses are uncontrolled by timing hardware logic components.

(14) Other embodiments may include a distributed processor, comprising: a substrate; a memory array disposed on the substrate, the memory array including a plurality of discrete memory banks; and a processing array disposed on the substrate, the processing array including a plurality of processor subunits, each one of the processor subunits being associated with a corresponding, dedicated one of the plurality of discrete memory banks; and a plurality of buses, each one of the plurality of buses connecting one of the plurality of processor subunits to at least another one of the plurality of processor subunits, wherein the plurality of processor subunits are configured to execute software that controls timing of data transfers across the plurality of buses to avoid colliding data transfers on at least one of the plurality of buses.

(15) Other embodiments may include a distributed processor on a memory chip, comprising: a substrate; a plurality of processor subunits disposed on the substrate, each processor subunit being configured to execute a series of instructions independent from other processor subunits, each series of instructions defining a series of tasks to be performed by a single processor subunit, a corresponding plurality of memory banks disposed on the substrate, each one of the plurality

processor subunits being connected to at least one dedicated memory bank not shared by any others of the plurality of processor subunits; and a plurality of buses, each of the plurality of buses connecting one of the plurality of processor subunits to at least one other of the plurality of processor subunits, wherein data transfers across at least one of the plurality of buses are predefined by the series of instructions included in a processor subunit connected to the at least one of the plurality of buses.

(16) Other embodiments may include a distributed processor on a memory chip, comprising: a plurality of processor subunits disposed on the memory chip; a plurality of memory banks disposed on the memory chip, wherein each one of the plurality of memory banks is configured to store data independent from data stored in other ones of the plurality of memory banks, and wherein each one of the plurality of processor subunits is connected to at least one dedicated memory bank from among the plurality of memory banks; and a plurality of buses, wherein each one of the plurality of buses connects one of the plurality of processor subunits to one or more corresponding, dedicated memory banks from among the plurality of memory banks, wherein data transfers across a particular one of the plurality of buses are controlled by a corresponding processor subunit connected to the particular one of the plurality of buses.

(17) Other embodiments may include a distributed processor on a memory chip, comprising: a plurality of processor subunits disposed on the memory chip; a plurality of memory banks disposed on the memory chip, wherein each one of the plurality of processor subunits is connected to at least one dedicated memory bank from among the plurality of memory banks, and wherein each memory bank of the plurality of memory banks is configured to store data independent from data stored in other ones of the plurality of memory banks, and wherein at least some of the data stored in one particular memory bank from among the plurality of memory banks comprises a duplicate of data stored in at least another one of the plurality of memory banks; and a plurality of buses, wherein each one of the plurality of buses connects one of the plurality of processor subunits to one or more corresponding, dedicated memory banks from among the plurality of memory banks, wherein data transfers across a particular one of the plurality of buses are controlled by a corresponding processor subunit connected to the particular one of the plurality of buses.

(18) Other embodiments may include a distributed processor on a memory chip, comprising: a plurality of processor subunits disposed on the memory chip; a plurality of memory banks disposed on the memory chip, wherein each one of the plurality of processor subunits is connected to at least one dedicated memory bank from among the plurality of memory banks, and wherein each memory bank of the plurality of memory banks is configured to store data independent from data stored in other ones of the plurality of memory banks, and wherein at least some of the data stored in one particular memory bank from among the plurality of memory banks comprises a duplicate of data stored in at least another one of the plurality of memory banks; and a plurality of buses, wherein each one of the plurality of buses connects one of the plurality of processor subunits to one or more corresponding, dedicated memory banks from among the plurality of memory banks, wherein data transfers across a particular one of the plurality of buses are controlled by a corresponding processor subunit connected to the particular one of the plurality of buses.

(19) Other embodiments may include a non-transitory computer-readable medium storing instructions for compiling a series of instructions for execution on a memory chip comprising a plurality of processor subunits and a plurality of memory banks, wherein each processor subunit from among the plurality of processor subunits is connected to at least one corresponding, dedicated memory bank from among the plurality of memory banks, the instructions causing at least one processor to: divide the series of instructions into a plurality of groups of sub-series instructions, the division comprising: assigning tasks associated with the series of instructions to different ones of the processor subunits, wherein the processor subunits are spatially distributed among the plurality of memory banks disposed on the memory chip; generating tasks to transfer data between pairs of the processor subunits of the memory chip, each pair of processor subunits

being connected by a bus, and grouping the assigned and generated tasks into the plurality of groups of sub-series instructions, wherein each of the plurality of groups of sub-series instructions corresponds to a different one of the plurality of processor sub-units; generate machine code corresponding to each of the plurality of groups of sub-series instructions; and assign the generated machine code corresponding to each of the plurality of groups of sub-series instructions to a corresponding one of the plurality of processor subunits in accordance with the division.

(20) Other embodiments may include a memory chip, comprising: a plurality of memory banks, each memory bank having a bank row decoder, a bank column decoder, and a plurality of memory sub-banks, each memory sub-bank having a sub-bank row decoder and a sub-bank column decoder for allowing reads and writes to locations on the memory sub-bank, each memory sub-bank comprising: a plurality of memory mats, each memory mat having a plurality of memory cells, wherein the sub-bank row decoders and the sub-bank column decoders are connected to the bank row decoder and the bank column decoder.

(21) Other embodiments may include a memory chip, comprising: a plurality of memory banks, each memory bank having a bank controller and a plurality of memory sub-banks, each memory sub-bank having a sub-bank row decoder and a sub-bank column decoder for allowing reads and writes to locations on the memory sub-bank, each memory sub-bank comprising: a plurality of memory mats, each memory mat having a plurality of memory cells, wherein the sub-bank row decoders and the sub-bank column decoders process read and write requests from the bank controller.

(22) Other embodiments may include a memory chip, comprising: a plurality of memory banks, each memory bank having a having a bank controller for processing reads and writes to locations on the memory bank, each memory bank comprising: a plurality of memory mats, each memory mat having a plurality of memory cells and having a mat row decoder and a mat column decoder, wherein the mat row decoders and the mat column decoders process read and write requests from the sub-bank controller.

(23) Other embodiments may include a memory chip, comprising: a plurality of memory banks, each memory bank having a bank controller, a row decoder, and a column decoder for allowing reads and writes to locations on the memory bank; and a plurality of buses connecting each controller of the plurality of bank controllers to at least one other controller of the plurality of bank controllers.

(24) One aspect of the present disclosure is directed to a memory device including a substrate; a plurality of memory banks on the substrate; a plurality of primary logic blocks on the substrate, each of the plurality of primary logic blocks being connected to at least one of the plurality of memory banks; a plurality of redundant blocks on the substrate, each of the plurality of redundant blocks being connected to at least one of the memory banks, each of the plurality of redundant blocks replicating at least one of the plurality of primary logic blocks; and a plurality of configuration switches on the substrate, each one of the plurality of the configuration switches being connected to at least one of the plurality of primary logic blocks or to at least one of the plurality of redundant blocks. In the memory device, upon detection of a fault associated with one of the plurality of primary logic blocks: a first configuration switch of the plurality of configuration switches may be configured to disable the one of the plurality of primary logic blocks, and a second configuration switch of the plurality of configuration switches may be configured to enable one of the plurality of redundant blocks that replicates the one of the plurality of primary logic blocks.

(25) Another aspect of the present disclosure is directed to a distributed processor on a memory chip including a substrate; an address manager on the substrate; a plurality of primary logic blocks on the substrate, each of the plurality of primary logic blocks being connected to at least one of the plurality of memory banks; a plurality of redundant blocks on the substrate, each of the plurality of redundant blocks being connected to at least one of the plurality of memory banks, each of the plurality of redundant blocks replicating at least one of the plurality of primary logic blocks; and a

bus on the substrate connected to each of the plurality of primary logic blocks, each of the plurality of redundant blocks, and the address manager. In the processor may assign running ID numbers to blocks in the plurality of primary logic blocks that pass a testing protocol; assign illegal ID numbers to blocks in the plurality of primary logic blocks that do not pass the testing protocol; and assign running ID numbers to blocks in the plurality of redundant blocks that pass the testing protocol.

(26) Yet another aspect of the present disclosure is directed to a method for configuring a distributed processor on a memory chip. The method may include: testing each one of a plurality of primary logic blocks on the substrate of the memory chip for at least one circuit functionality, identifying at least one faulty logic block in the plurality of primary logic blocks based on the testing results, the at least one faulty logic block being connected to at least one memory bank disposed on the substrate of the memory chip; testing at least one redundant block on the substrate of the memory chip for the at least one circuit functionality, the at least one redundant block replicating the at least one faulty logic block and being connected to the at least one memory bank; disabling the at least one faulty logic block by applying an external signal to a deactivation switch, the deactivation switch being connected with the at least one faulty logic block and being disposed on the substrate of the memory chip; and enabling the at least one redundant block by applying the external signal to an activation switch, the activation switch being connected with the at least one redundant block and being disposed on the substrate of the memory chip.

(27) Another aspect of the present disclosure is directed to a method for configuring a distributed processor on a memory chip. The method may include enabling a plurality of primary logic blocks and a plurality of redundant blocks on the substrate of the memory; testing each one of the plurality of primary logic blocks on the substrate of the memory chip for at least one circuit functionality; identifying at least one faulty logic block in the plurality of primary logic blocks based on the testing results, the at least one faulty logic block being connected to at least one memory bank disposed on the substrate of the memory chip; testing at least one redundant block on the substrate of the memory chip for the at least one circuit functionality, the at least one redundant block replicating the at least one faulty logic block and being connected to the at least one memory bank; and disabling at least one redundant block by applying the external signal to an activation switch, the activation switch being connected with the at least one redundant block and being disposed on the substrate of the memory chip.

(28) One aspect of the present disclosure is directed to a processing device. The processing device may include a substrate; a plurality of memory banks on the substrate; a memory controller on the substrate connected to each one of the plurality of memory banks; and a plurality of processing units on the substrate, each one of the plurality of processing units being connected to the memory controller, the plurality of processing units including a configuration manager. In the processing device, the configuration manager is configured to receive a first indication of a task to be performed, the task requiring at least one computation; signal at least one selected processing unit from the plurality of processing units based upon a capability of the selected processing unit for performing the at least one computation; and transmitting a second indication to the at least one selected processing unit, and the memory controller is configured to route data from at least two memory banks to the at least one selected processing unit using at least one communication line, the at least one communication line being connected to the at least two memory banks and the at least one selected processing unit via the memory controller.

(29) Another aspect of the present disclosure is directed to a method performed for operating a distributed memory device. The method may include: compiling, by a compiler, a task for the distributed memory device, the task requiring at least one computation, the compiling may include determining a number of words that are required simultaneously to perform the task, and providing instructions for writing words that need to be accessed simultaneously in a plurality of memory banks disposed on the substrate when a number a number of words that can be accessed

simultaneously from one of the plurality of memory banks is lower than the number of words that are required simultaneously; receiving, by a configuration manager disposed on the substrate, an indication to perform the task; and in response to receiving the indication, configuring a memory controller disposed in the substrate to: within a first line access cycle: access at least one first word from a first memory bank from the plurality of memory banks using a first memory line, send the at least one first word to at least one processing unit, and open a first memory line in the second memory bank to access a second address from the second memory bank from the plurality of memory banks, and within a second line access cycle: access at least one second word from the second memory bank using the first memory line, send the at least one second word to at least one processing unit, and access a third address from the first memory bank using a second memory line in the first bank.

(30) Yet another aspect of the present disclosure is directed to anon-transitory computer-readable medium that stores instructions that, when executed by at least one processor, cause the at least one processor to determine a number of words that are required simultaneously to perform a task, the task requiring at least one computation; write words that need to be accessed simultaneously in a plurality of memory banks disposed on the substrate when a number a number of words that can be accessed simultaneously from one of the plurality of memory banks is lower than the number of words that are required simultaneously; transmit an indication to perform the task to a configuration manager disposed on the substrate; and transmit instructions to configure a memory controller disposed on the substrate to, within a first line access cycle: access at least one first word from a first memory bank from the plurality of memory banks using a first memory line, send the at least one first word to at least one processing unit, and open a first memory line in the second memory bank to access a second address from the second memory bank from the plurality of memory banks, and within a second line access cycle: access at least one second word from the second memory bank using the first memory line, send the at least one second word to at least one processing unit, and access a third address from the first memory bank using a second memory line in the first bank.

(31) Consistent with other disclosed embodiments, non-transitory computer-readable storage media may store program instructions, which are executed by at least one processing device and perform any of the methods described herein.

(32) The foregoing general description and the following detailed description are exemplary and explanatory only and are not restrictive of the claims.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The accompanying drawings, which are incorporated in and constitute a part of this disclosure, illustrate various disclosed embodiments. In the drawings:

(2) FIG. 1 is a diagrammatic representation of a central processing unit (CPU).

(3) FIG. 2 is a diagrammatic representation of a graphics processing unit (GPU).

(4) FIG. 3A is a diagrammatic representation of an embodiment of an exemplary hardware chip consistent with the disclosed embodiments.

(5) FIG. 3B is a diagrammatic representation of another embodiment of an exemplary hardware chip consistent with the disclosed embodiments.

(6) FIG. 4 is a diagrammatic representation of a generic command executed by an exemplary hardware chip consistent with the disclosed embodiments.

(7) FIG. 5 is a diagrammatic representation of a specialized command executed by an exemplary hardware chip consistent with the disclosed embodiments.

(8) FIG. 6 is a diagrammatic representation of a processing group for use in an exemplary hardware

chip consistent with the disclosed embodiments.

(9) FIG. 7A is a diagrammatic representation of a rectangular array of processing groups consistent with the disclosed embodiments.

(10) FIG. 7B is a diagrammatic representation of an elliptical array of processing groups consistent with the disclosed embodiments.

(11) FIG. 7C is a diagrammatic representation an array of hardware chips consistent with the disclosed embodiments.

(12) FIG. 7D is a diagrammatic representation another array of hardware chips consistent with the disclosed embodiments.

(13) FIG. 8 is a flowchart depicting an exemplary method for compiling a series of instructions for execution on an exemplary hardware chip consistent with the disclosed embodiments.

(14) FIG. 9 is a diagrammatic representation of a memory bank.

(15) FIG. 10 is a diagrammatic representation of a memory bank.

(16) FIG. 11 is a diagrammatic representation of an embodiment of an exemplary memory bank with sub-bank controls consistent with the disclosed embodiments.

(17) FIG. 12 is a diagrammatic representation of another embodiment of an exemplary memory bank with sub-bank controls consistent with the disclosed embodiments.

(18) FIG. 13 is a block diagram of an exemplary memory chip, consistent with disclosed embodiments.

(19) FIG. 14 is a block diagram of an exemplary redundant logic block set, consistent with disclosed embodiments.

(20) FIG. 15 is a block diagram for an exemplary logic block, consistent with disclosed embodiments.

(21) FIG. 16 are block diagrams of exemplary logic blocks connected with a bus, consistent with disclosed embodiments.

(22) FIG. 17 is a block diagram for exemplary logic blocks connected in series, consistent with disclosed embodiments.

(23) FIG. 18 is a block diagram of exemplary logic blocks connected in a two-dimension array, consistent with disclosed embodiments.

(24) FIG. 19 is a block diagram for exemplary logic blocks in a complex connection, consistent with disclosed embodiments.

(25) FIG. 20 is an exemplary flow chart illustrating a redundant block enabling process, consistent with disclosed embodiments.

(26) FIG. 21 is an exemplary flow chart illustrating an address assignment process, consistent with disclosed embodiments.

(27) FIG. 22 provides block diagrams for exemplary processing devices, consistent with disclosed embodiments.

(28) FIG. 23 is a block diagram of an exemplary processing device, consistent with disclosed embodiments.

(29) FIG. 24 includes exemplary memory configuration diagrams, consistent with disclosed embodiments.

(30) FIG. 25 is an exemplary flowchart illustrating a memory configuration process, consistent with disclosed embodiments.

(31) FIG. 26 is an exemplary flowchart illustrating a memory read process, consistent with disclosed embodiments.

(32) FIG. 27 is an exemplary flowchart illustrating a process execution, consistent with disclosed embodiments.

DETAILED DESCRIPTION

(33) The following detailed description refers to the accompanying drawings. Wherever convenient, the same reference numbers are used in the drawings and the following description to

refer to the same or similar parts. While several illustrative embodiments are described herein, modifications, adaptations and other implementations are possible. For example, substitutions, additions or modifications may be made to the components illustrated in the drawings, and the illustrative methods described herein may be modified by substituting, reordering, removing, or adding steps to the disclosed methods. Accordingly, the following detailed description is not limited to the disclosed embodiments and examples. Instead, the proper scope is defined by the appended claims.

(34) Processor Architecture

(35) As used throughout this disclosure, the term “hardware chip” refers to a semiconductor wafer (such as silicon or the like) on which one or more circuit elements (such as transistors, capacitors, resistors, and/or the like) are formed. The circuit elements may form processing elements or memory elements. A “processing element” refers to one or more circuit elements that, together, perform at least one logic function (such as an arithmetic function, a logic gate, other Boolean operations, or the like). A processing element may be a general-purpose processing element (such as a configurable plurality of transistors) or a special-purpose processing element (such as a particular logic gate or a plurality of circuit elements designed to perform a particular logic function). A “memory element” refers to one or more circuit elements that can be used to store data. A “memory element” may also be referred to as a “memory cell.” A memory element may be dynamic (such that electrical refreshes are required to maintain the data store), static (such that data persists for at least some time after power loss), or non-volatile memories.

(36) Processing elements may be joined to form processor subunits. A “processor subunit” may thus comprise a smallest grouping of processing elements that may execute at least one task or instructions (e.g., of a processor instruction set). For example, a subunit may comprise one or more general-purpose processing elements configured to execute instructions together, one or more general-purpose processing elements paired with one or more special-purpose processing elements configured to execute instructions in a complementary fashion, or the like. The processor subunits may be arranged on a substrate (e.g., a wafer) in an array. Although the “array” may comprise a rectangular shape, any arrangement of the subunits in the array may be formed on the substrate.

(37) Memory elements may be joined to form memory banks. For example, a memory bank may comprise one or more lines of memory elements linked along at least one wire (or other conductive connection). Furthermore, the memory elements may be linked along at least one additional wire in another direction. For example, the memory elements may be arranged along wordlines and bitlines, as explained below. Although the memory bank may comprise lines, any arrangement of the elements in the bank may be used to form the bank on the substrate. Moreover, one or more banks may be electrically joined to at least one memory controller to form a memory array. Although the memory array may comprise a rectangular arrangement of the banks, any arrangement of the banks in the array may be formed on the substrate.

(38) As further used throughout this disclosure, a “bus” refers to any communicative connection between elements of a substrate. For example, a wire or a line (forming an electrical connection), an optical fiber (forming an optical connection), or any other connection conducting communications between components may be referred to as a “bus.”

(39) Conventional processors pair general-purpose logic circuits with shared memories. The shared memories may store both instruction sets for execution by the logic circuits as well as data used for and resulting from execution of the instruction sets. As described below, some conventional processors use a caching system to reduce delays in performing pulls from the shared memory; however, conventional caching systems remain shared. Conventional processors include central processing units (CPUs), graphics processing units (GPUs), various application-specific integrated circuits (ASICs), or the like. FIG. 1 shows an example of a CPU, and FIG. 2 shows an example of a GPU.

(40) As shown in FIG. 1, a CPU **100** may comprise a processing unit **110** that includes one or more

processor subunits, such as processor subunit **120a** and processor subunit **120b**. Although not depicted in FIG. 1, each processor subunit may comprise a plurality of processing elements. Moreover, the processing unit **110** may include one or more levels of on-chip cache. Such cache elements are generally formed on the same semiconductor die as processing unit **110** rather than being connected to processor subunits **120a** and **120b** via one or more buses formed in the substrate containing processor subunits **120a** and **120b** and the cache elements. An arrangement directly on the same die, rather than being connected via buses, is common for both first-level (L1) and second-level (L2) caches in conventional processors. Alternatively, in older processors, L2 caches were shared amongst processor subunits using back-side buses between the subunits and the L2 caches. Back-side buses are generally larger than front-side buses, described below. Accordingly, because cache is to be shared with all processor subunits on the die, cache **130** may be formed on the same die as processor subunits **120a** and **120b** or communicatively coupled to processor subunits **120a** and **120b** via one or more back-side buses. In both embodiments without buses (e.g., cache is formed directly on-die) as well as embodiments using back-side buses, the caches are shared between processor subunits of the CPU.

(41) Moreover, processing unit **110** communicates with shared memory **140a** and memory **140b**. For example, memories **140a** and **140b** may represent memory banks of shared dynamic random access memory (DRAM). Although depicted with two banks, most conventional memory chips include between eight and sixteen memory banks. Accordingly, processor subunits **120a** and **120b** may use shared memories **140a** and **140b** to store data that is then operated upon by processor subunits **120a** and **120b**. This arrangement, however, results in the buses between memories **140a** and **140b** and processing unit **110** acting as a bottleneck when the clock speeds of processing unit **110** exceed data transfer speeds of the buses. This is generally true for conventional processors, resulting in lower effective processing speeds than the stated processing speeds based on clock rate and number of transistors.

(42) As shown in FIG. 2, similar deficiencies persist in GPUs. A GPU **200** may comprise a processing unit **210** that includes one or more processor subunits (e.g., subunits **220a**, **220b**, **220c**, **220d**, **220e**, **220f**, **220g**, **220h**, **220i**, **220j**, **220k**, **220l**, **220m**, **220n**, **220o**, and **220p**). Moreover, the processing unit **210** may include one or more levels of on-chip cache and/or register files. Such cache elements are generally formed on the same semiconductor die as processing unit **210**. Indeed, in the example of FIG. 2, cache **210** is formed on the same die as processing unit **210** and shared amongst all of the processor subunits, while caches **230a**, **230b**, **230c**, and **230d** are formed on a subset of the processor subunits, respectively, and dedicated thereto.

(43) Moreover, processing unit **210** communicates with shared memories **250a**, **250b**, **250c**, and **250d**. For example, memories **250a**, **250b**, **250c**, and **250d** may represent memory banks of shared DRAM. Accordingly, the processor subunits of processing unit **210** may use shared memories **250a**, **250b**, **250c**, and **250d** to store data that is then operated upon by the processor subunits. This arrangement, however, results in the buses between memories **250a**, **250b**, **250c**, and **250d** and processing unit **210** acting as a bottleneck, similar to the bottleneck described above for CPUs.

(44) Overview of Disclosed Hardware Chips

(45) FIG. 3A is a diagrammatic representation of an embodiment depicting an exemplary hardware chip **300**. Hardware chip **300** may comprise a distributed processor designed to mitigate the bottlenecks described above for CPUs, GPUs, and other conventional processors. A distributed processor may include a plurality of processor subunits distributed spatially on a single substrate. Moreover, as explained above, in distributed processors of the present disclosure, corresponding memory banks are also spatially distributed on the substrate. In some embodiments, a distributed processor may be associated with a set of instructions, and each one of the processor subunits of the distributed processor may be responsible for performing one or more tasks included in the set of instructions.

(46) As depicted in FIG. 3A, hardware chip **300** may comprise a plurality of processor subunits,

e.g., logic and control subunits **320a**, **320b**, **320c**, **320d**, **320e**, **320f**, **320g**, and **320h**. As further depicted in FIG. 3A, each processor subunit may have a dedicated memory instance. For example, logic and control subunit **320a** is operably connected to dedicated memory instance **330a**, logic and control subunit **320b** is operably connected to dedicated memory instance **330b**, logic and control subunit **320c** is operably connected to dedicated memory instance **330c**, logic and control subunit **320d** is operably connected to dedicated memory instance **330d**, logic and control subunit **320e** is operably connected to dedicated memory instance **330e**, logic and control subunit **320f** is operably connected to dedicated memory instance **330f**, logic and control subunit **320g** is operably connected to dedicated memory instance **330g**, and logic and control subunit **320h** is operably connected to dedicated memory instance **330h**.

(47) Although FIG. 3A depicts each memory instance as a single memory bank, hardware chip **300** may include two or more memory banks as a dedicated memory instance for a processor subunit on hardware chip **300**. Furthermore, although FIG. 3A depicts each processor subunit as comprising both a logic component and a control for the dedicated memory bank(s), hardware chip **300** may use controls for the memory banks that are separate, at least in part, from the logic components. Moreover, as depicted in FIG. 3A, two or more processor subunits and their corresponding memory banks may be grouped, e.g., into processing groups **310a**, **310b**, **310c**, and **310d**. A “processing group” may represent a spatial distinction on a substrate on which hardware chip **300** is formed. Accordingly, a processing group may include further controls for the memory banks in the group, e.g., controls **340a**, **340b**, **340c**, and **340d**. Additionally or alternatively, a “processing group” may represent a logical grouping for the purposes of compiling code for execution on hardware chip **300**. Accordingly, a compiler for hardware chip **300** (further described below) may divide an overall set of instructions between the processing groups on hardware chip **300**.

(48) Furthermore, host **350** may provide instructions, data, and other input to hardware chip **300** and read output from the same. Accordingly, a set of instructions may be executed entirely on a single die, e.g., the die hosting hardware chip **300**. Indeed, the only communications off-die may include the loading of instructions to hardware chip **300**, any input sent to hardware chip **300**, and any output read from hardware chip **300**. Accordingly, all calculations and memory operations may be performed on-die (on hardware chip **300**) because the processor subunits of hardware chip **300** communicate with dedicated memory banks of hardware chip **300**.

(49) FIG. 3B is a diagrammatic representation of an embodiment depicting another exemplary hardware chip **300'**. Although depicted as an alternative to hardware chip **300**, the architecture depicted in FIG. 3B may be combined, at least in part, with the architecture depicted in FIG. 3A.

(50) As depicted in FIG. 3B, hardware chip **300'** may comprise a plurality of processor subunits, e.g., processor subunits **350a**, **350b**, **350c**, and **350d**. As further depicted in FIG. 3B, each processor subunit may have a plurality of dedicated memory instances. For example, processor subunit **350a** is operably connected to dedicated memory instances **330a** and **330b**, processor subunit **350b** is operably connected to dedicated memory instances **330c** and **330d**, processor subunit **350c** is operably connected to dedicated memory instances **330e** and **330f**, and processor subunit **350d** is operably connected to dedicated memory instances **330g** and **330h**. Moreover, as depicted in FIG. 3B, the processor subunits and their corresponding memory banks may be grouped, e.g., into processing groups **310a**, **310b**, **310c**, and **310d**. As explained above, a “processing group” may represent a spatial distinction on a substrate on which hardware chip **300'** is formed and/or a logical grouping for the purposes of compiling code for execution on hardware chip **300'**.

(51) As further depicted in FIG. 3B, the processor subunits may communicate with each other via buses. For example, as shown in FIG. 3B, processor subunit **350a** may communicate with processor subunit **350b** via bus **360a**, with processor subunit **350c** via bus **360c**, and with processor subunit **350d** via bus **360f**. Similarly, processor subunit **350b** may communicate with processor subunit **350a** via bus **360a** (as described above), with processor subunit **350c** via bus **360e**, and with processor subunit **350d** via bus **360d**. In addition, processor subunit **350c** may communicate

with processor subunit **350a** via bus **360c** (as described above), with processor subunit **350b** via bus **360e** (as described above), and with processor subunit **350d** via bus **360b**. Accordingly, processor subunit **350d** may communicate with processor subunit **350a** via bus **360f** (as described above), with processor subunit **350b** via bus **360d** (as described above), and with processor subunit **350c** via bus **360b** (as described above). One of ordinary skill will understand that fewer buses than depicted in FIG. 3B may be used. For example, bus **360e** may be eliminated such that communications between processor subunit **350b** and **350c** pass through processor subunit **350a** and/or **350d**. Similarly, bus **360f** may be eliminated such that communications between processor subunit **350a** and processor subunit **350d** pass through processor subunit **350b** or **350c**.

(52) Moreover, one of ordinary skill will understand that architectures other than those depicted in FIGS. 3A and 3B may be used. For example, an array of processing groups, each with a single processor subunit and memory instance, may be arranged on a substrate. Processor subunits may additionally or alternatively form part of controllers for corresponding dedicated memory banks, part of controllers for memory mats of corresponding dedicated memory, or the like.

(53) In view of the architecture described above, hardware chips **300** and **300'** may provide significant increases in efficiency for memory-intensive tasks as compared with traditional architectures. For example, database operations and artificial intelligence algorithms (such as neural networks) are examples of memory-intensive tasks for which traditional architectures are less efficient than hardware chips **300** and **300'**. Accordingly, hardware chips **300** and **300'** may be referred to as database accelerator processors and/or artificial intelligence accelerator processors.

(54) Configuring the Disclosed Hardware Chips

(55) The hardware chip architecture described above may be configured for execution of code. For example, each processor subunit may individually execute code (defining a set of instructions) apart from other processor subunits in the hardware chip. Accordingly, rather than relying on an operating system to manage multithreading or using multitasking (which is concurrency rather than parallelism), hardware chips of the present disclosure may allow for processor subunits to operate fully in parallel.

(56) In addition to a fully parallel implementation described above, at least some of the instructions assigned to each processor subunit may be overlapping. For example, a plurality of processor subunits on a distributed processor may execute overlapping instructions as, for example, an implementation of an operating system or other management software, while executing non-overlapping instructions in order to perform parallel tasks within the context of the operating system or other management software.

(57) FIG. 4 depicts an exemplary process **400** for executing a generic command with processing group **410**. For example, processing group **410** may comprise a portion of a hardware chip of the present disclosure, e.g., hardware chip **300**, hardware chip **300'**, or the like.

(58) As depicted in FIG. 4, a command may be sent to processor subunit **430**, which is paired with dedicated memory instance **420**. An external host (e.g., host **350**) may send the command to processing group **410** for execution. Alternatively, host **350** may have sent an instruction set including the command for storage in memory instance **420** such that processor subunit **430** may retrieve the command from memory instance **420** and execute the retrieved command. Accordingly, the command may be executed by processing element **440**, which is a generic processing element configurable to execute the received command. Moreover, processing group **410** may include a control **460** for memory instance **420**. As depicted in FIG. 4, control **460** may perform any reads and/or writes to memory instance **420** required by processing element **440** when executing the received command. After execution of the command, processing group **410** may output the result of the command, e.g., to the external host or to a different processing group on the same hardware chip.

(59) In some embodiments, as depicted in FIG. 4, processor subunit **430** may further include an address generator **450**. An “address generator” may comprise a plurality of processing elements

that are configured to determine addresses in one or more memory banks for performing reads and writes and may also perform operations on the data located at the determined addresses (e.g., addition, subtraction, multiplication, or the like). For example, address generator **450** may determine addresses for any reads or writes to memory. In one example, address generator **450** may increase efficiency by overwriting a read value with a new value determined based on the command when the read value is no longer needed. Additionally or alternatively, address generator **450** may select available addresses for storage of results from execution of the command. This may allow for scheduling of result read-off for a later clock cycle, when it is more convenient for the external host. In another example, address generator **450** may determine addresses to read from and write to during a multi-cycle calculation, such as a vector or matrix multiply-accumulate calculation. Accordingly, address generator **450** may maintain or calculate memory addresses for reading data and writing intermediate results of the multi-cycle calculation such that processor subunit **430** may continue processing without having to store these memory addresses.

(60) FIG. 5 depicts an exemplary process **500** for executing a specialized command with processing group **510**. For example, processing group **510** may comprise a portion of a hardware chip of the present disclosure, e.g., hardware chip **300**, hardware chip **300'**, or the like.

(61) As depicted in FIG. 5, a specialized command (e.g., a multiply-accumulate command) may be sent to processing element **530**, which is paired with dedicated memory instance **520**. An external host (e.g., host **350**) may send the command to processing element **530** for execution. Accordingly, the command may be executed at a given signal from the host by processing element **530**, a specialized processing element configurable to execute particular commands (including the received command). Alternatively, processing element **530** may retrieve the command from memory instance **520** for execution. Thus, in the example of FIG. 5, processing element **530** is a multiply-accumulate (MAC) circuit configured to execute MAC commands received from the external host or retrieved from memory instance **520**. After execution of the command, processing group **410** may output the result of the command, e.g., to the external host or to a different processing group on the same hardware chip. Although depicted with a single command and a single result, a plurality of commands may be received or retrieved and executed, and a plurality of results may be combined on processing group **510** before output.

(62) Although depicted as a MAC circuit in FIG. 5, additional or alternative specialized circuits may be included in processing group **510**. For example, a MAX-read command (which returns the max value of a vector) a MAX0-read command (a common function also termed a rectifier, which returns the entire vector but also does MAX with 0), or the like may be implemented.

(63) Although depicted separately, the generalized processing group **410** of FIG. 4 and the specialized processing group **510** of FIG. 5 may be combined. For example, a generic processor subunit may be coupled to one or more specialized processor subunits to form a processor subunit. Accordingly, the generic processor subunit may be used for all instructions not executable by the one or more specialized processor subunits.

(64) One of ordinary skill will understand that neural network implementation and other memory-intensive tasks may be handled with specialized logic circuits. For example, database queries, packet inspection, string comparison, and other functions may increase in efficiency if executed by the hardware chips described herein.

(65) A Memory-Based Architecture for Distributed Processing

(66) On hardware chips consistent with the present disclosure, dedicated buses may transfer data between processor subunits on the chip and/or between the processor subunits and their corresponding dedicated memory banks. The use of dedicated buses may reduce arbitration costs because competing requests are either not possible or easily avoided using software rather than hardware.

(67) FIG. 6 schematically depicts a diagrammatic representation of a processing group **600**.

Processing group **600** may be for use in a hardware chip, e.g., hardware chip **300**, hardware chip

300', or the like. Processor subunit **610** may be connected via buses **630** to memory **620**. Memory **620** may comprise a Randomly Accessible Memory (RAM) element that stores data and code for execution by processor subunit **610**. In some embodiments, memory **620** may be an N-Way memory (wherein N is a number equal to or larger than 1 that implies the number of segments in an interleaved memory **620**). Because processor subunit **610** is coupled to memory **620** dedicated to processor subunit **610** via bus **630**, N may be kept relatively small without compromising the execution performance. This represents an improvement over conventional multiway register files or caches where a lower N generally results in lower execution performance, and a higher N generally results in large area and power loss.

(68) The size of memory **620**, the number of ways, and the width of bus **630** may be adjusted to meet the requirements of tasks and application implementations of a system using processing group **600** according to, for instance, the size of data involved in the task or tasks. Memory element **620** may comprise one or more types of memory known in the art, e.g., volatile memory (such as RAM, DRAM, SRAM, phase-change RAM (PRAM), magnetoresistive RAM (MRAM), resistive RAM (ReRAM), or the like) or non-volatile memory (such as flash or ROM). According to some embodiments, a portion of memory element **620** may comprise a first memory type, while another portion may comprise another memory type. For instance, the code region of a memory element **620** may comprise a ROM element, while a data region of the memory element **620** may comprise a DRAM element. Another example for such partitioning is storing the weights of a neural network in flash while storing the data for calculation in DRAM.

(69) Processor subunit **610** comprises a processing element **640** that may comprise a processor. The processor can be pipelined or not pipelined, a customized Reduced Instruction Set Computing (RISC) element or other processing scheme, implemented on any commercial Integrated Circuit (IC) known in the art (such as ARM, ARC, RISC-V, etc.), as appreciated by one of ordinary skill. Processing element **640** may comprise a controller that, in some embodiments, includes an Arithmetic Logic Unit (ALU) or other controller.

(70) According to some embodiments, processing element **640**, which executes received or stored code, may comprise a generic processing element and, therefore, be flexible and capable of performing a wide variety of processing operations. Non-dedicated circuitry typically consumes more power than specific-operation-dedicated circuitry when comparing the power consumed during performance for a specific operation. Therefore, when performing specific complex arithmetic calculations, processing element **640** may consume more power and perform less efficiently than dedicated hardware. Therefore, according to some embodiments, a controller of processing element **640** may be designed to perform specific operations (e.g., addition or “move” operations).

(71) In one example, the specific operations may be performed by one or more accelerators **650**. Each accelerator may be dedicated and programmed to perform a specific calculation (such as multiplication, floating point vector operations, or the like). By using accelerator(s), the average power consumed per calculation per processor subunit may be lowered, and the calculation throughput henceforth increases. Accelerator(s) **650** may be chosen according to an application that the system is designed to implement (e.g., execution of neural networks, execution of database queries, or the like). Accelerator(s) **650** may be configured by processing element **640** and may operate in tandem therewith for lowering power consumption and accelerating calculations and computations. The accelerators may additionally or alternatively be used to transfer data between memory and MUXs/DEMUXs/input/output ports (e.g., MUX **650** and DEMUX **660**) of processing group **600**, such as a smart DMA (direct memory access) peripheral.

(72) Accelerator(s) **650** may be configured to perform a variety of functions. For instance, one accelerator may be configured to perform 16-bit floating point calculation or 8-bit integer calculations, which are often used in neural networks. Another example of an accelerator function is a 32-bit floating point calculation, which is often used during a training stage of a neural

network. Yet another example of an accelerator function is query processing, such as that used in databases. In some embodiments, accelerator(s) **650** may comprise specialized processing elements to perform these functions and/or may be configured according to configuration data, stored on the memory element **620**, such that it may be modified.

(73) Accelerator(s) **650** may additionally or alternatively implement a configurable scripted list of memory movements to time movements of data to/from memory **620** or to/from other accelerators and/or inputs/outputs. Accordingly, as explained further below, all the data movement inside the hardware chip using processing group **600** may use software synchronization rather than hardware synchronization. For example, an accelerator in one processing group (e.g., group **600**) may transfer data from its input to its accelerator every tenth cycle and then output data at the next cycle, thereby letting the information flow from the memory of the processing group to another one.

(74) As further depicted in FIG. **6**, in some embodiments, processing group **600** may further comprise at least one input multiplexer (MUX) **660** connected to its input port and at least one output DEMUX **670** connected to its output port. These MUXs/DEMUXs may be controlled by control signals (not shown) from processing element **640** and/or from one of accelerator(s) **650**, determined according to a current instruction being carried out by processing element **640** and/or the operation executed by an accelerator of accelerator(s) **650**. In some scenarios, processing group **600** may be required (according to a predefined instruction from its code memory) to transfer data from its input port to its output port. Accordingly, one or more of the input MUXs (e.g., MUX **660**) may be directly connected via one or more buses to an output DEMUX (e.g., DEMUX **670**), in addition to each of the DEMUXs/MUXs being connected to processing element **640** and accelerator(s) **650**.

(75) The processing group **600** of FIG. **6** may be arrayed to form a distributed processor, for example, as depicted in FIG. **7A**. The processing groups may be disposed on substrate **710** to form an array. In some embodiments, substrate **710** may comprise a semiconductor substrate, such as silicon. Additionally or alternatively, substrate **710** may comprise a circuit board, such as a flexible circuit board.

(76) As depicted in FIG. **7A**, substrate **710** may include, disposed thereon, a plurality of processing groups, such as processing group **600**. Accordingly, substrate **710** includes a memory array that includes a plurality of banks, such as banks **720a**, **720b**, **720c**, **720d**, **720e**, **720f**, **720g**, and **720h**. Furthermore, substrate **710** includes a processing array that may include a plurality of processor subunits, such as subunits **730a**, **730b**, **730c**, **730d**, **730e**, **730f**, **730g**, and **730h**.

(77) Furthermore, as explained above, each processing group may include a processor subunit and one or more corresponding memory banks dedicated to the processor subunit. Accordingly, as depicted in FIG. **7A**, each subunit is associated with a corresponding, dedicated memory bank, e.g.: Processor subunit **730a** is associated with memory bank **720a**, processor subunit **730b** is associated with memory bank **720b**, processor subunit **730c** is associated with memory bank **720c**, processor subunit **730d** is associated with memory bank **720d**, processor subunit **730e** is associated with memory bank **720e**, processor subunit **730f** is associated with memory bank **720f**, processor subunit **730g** is associated with memory bank **720g**, processor subunit **730h** is associated with memory bank **720h**.

(78) To allow each processor subunit to communicate with its corresponding, dedicated memory bank(s), substrate **710** may include a first plurality of buses connecting one of the processor subunits to its corresponding, dedicated memory bank(s). Accordingly, bus **740a** connects processor subunit **730a** to memory bank **720a**, bus **740b** connects processor subunit **730b** to memory bank **720b**, bus **740c** connects processor subunit **730c** to memory bank **720c**, bus **740d** connects processor subunit **730d** to memory bank **720d**, bus **740e** connects processor subunit **730e** to memory bank **720e**, bus **740f** connects processor subunit **730f** to memory bank **720f**, bus **740g** connects processor subunit **730g** to memory bank **720g**, and bus **740h** connects processor subunit

730h to memory bank 720h. Moreover, to allow each processor subunit to communicate with other processor subunits, substrate 710 may include a second plurality of buses connecting one of the processor subunits to another of the processor subunits. In the example of FIG. 7A, bus 750a connects processor subunit 730a to processor subunit 750e, bus 750b connects processor subunit 730a to processor subunit 750b, bus 750c connects processor subunit 730b to processor subunit 750f, bus 750d connects processor subunit 730b to processor subunit 750c, bus 750e connects processor subunit 730c to processor subunit 750g, bus 750f connects processor subunit 730c to processor subunit 750d, bus 750g connects processor subunit 730d to processor subunit 750h, bus 750h connects processor subunit 730h to processor subunit 750g, bus 750i connects processor subunit 730g to processor subunit 750g, and bus 750j connects processor subunit 730f to processor subunit 750e.

(79) Accordingly, in the example arrangement shown in FIG. 7A, the plurality of logic processor subunits is arranged in at least one row and at least one column. The second plurality of buses connect each processor subunit to at least one adjacent processor subunit in the same row and to at least one adjacent processor subunit in the same column. FIG. 7A may be referred to as a “partial file connection.”

(80) The arrangement shown in FIG. 7A may be modified to form a “full tile connection.” A full tile connection includes additional buses connecting diagonal processor subunits. For example, the second plurality of buses may include additional buses between processor subunit 730a and processor subunit 730f, between processor subunit 730b and processor subunit 730e, between processor subunit 730b and processor subunit 730g, between processor subunit 730c and processor subunit 730f, between processor subunit 730c and processor subunit 730h, and between processor subunit 730d and processor subunit 730g.

(81) A full tile connection may be used for convolution calculations, in which data and results stored in a near processor subunit are used. For example, during convolutional image processing, each processor subunit may receive a tile of the image (such as a pixel or a group of pixels). In order to calculate the convolution results, each processor subunit may acquire data from all eight adjacent processor subunits, each of which have received a corresponding tile. In a partial tile connection, the data from the diagonal adjacents may be passed through other adjacent processor subunits connected to the processor subunit. Accordingly, the distributed processor on a chip may be an artificial intelligence accelerator processor.

(82) In a specific example of a convolutional calculation, an $N \times M$ image may be divided across a plurality of processor subunits. Each processor subunit may perform a convolution with an $A \times B$ filter on its corresponding tile. To perform the filtering on one or more pixels on a boundary between tiles, each processor subunit may require data from neighboring processor subunits having tiles including pixels on the same boundary. Accordingly, the code generated for each processor subunit configures the subunit to calculate the convolutions and pull from one of the second plurality of buses whenever data is needed from an adjacent subunit. Corresponding commands to output data to the second plurality of buses are provided to the subunits to ensure proper timing of needed data transfers.

(83) The partial tile connection of FIG. 7A may be modified to be an N-partial tile connection. In this modification, the second plurality of buses may further connect each processor subunit to processor subunits within a threshold distance of the processor subunit (e.g., within n processor subunits) in the four directions along which the buses of FIG. 7A run (i.e., up, down, left, and right). A similar modification may be made to the full-tile connection (to result in an N-full tile connection) such that the second plurality of buses further connects each processor subunit to processor subunits within a threshold distance of the processor subunit (e.g., within n processor subunits) in the four directions along which the buses of FIG. 7A run in addition to the two diagonal directions.

(84) Other arrangements are possible. For example, in the arrangement shown in FIG. 7B, bus 750a

connects processor subunit **730a** to processor subunit **730d**, bus **750b** connects processor subunit **730a** to processor subunit **730b**, bus **750c** connects processor subunit **730b** to processor subunit **730c**, and bus **750d** connects processor subunit **730c** to processor subunit **730d**. Accordingly, in the example arrangement shown in FIG. 7B, the plurality of processor subunits is arranged in a star pattern. The second plurality of buses connect each processor subunit to at least one adjacent processor subunit within the star pattern.

(85) Further arrangements (not shown) are possible. For example, a neighbor connection arrangement may be used such that the plurality of processor subunits is arranged in one or more lines (e.g., similar to that depicted in FIG. 7A). In a neighbor connection arrangement, the second plurality of buses connect each processor subunit to a processor subunit to the left in the same line, to a processor subunit to the right in the same line, to the processor subunits both to the left and to the right in the same line, etc.

(86) In another example, an N-linear connection arrangement may be used. In an N-linear connection arrangement, the second plurality of buses connect each processor subunit to processor subunits within a threshold distance of the processor subunit (e.g., within n processor subunits). The N-linear connection arrangement may be used with the line array (described above), the rectangular array (depicted in FIG. 7A), the elliptical array (depicted in FIG. 7B), or any other geometrical array.

(87) In yet another example, an N-log connection arrangement may be used. In an N-log connection arrangement, the second plurality of buses connect each processor subunit to processor subunits within a threshold power of two distance of the processor subunit (e.g., within $2^{\sup.n}$ processor subunits). The N-log connection arrangement may be used with the line array (described above), the rectangular array (depicted in FIG. 7A), the elliptical array (depicted in FIG. 7B), or any other geometrical array.

(88) Any of the connection schemes described above may be combined for use in the same hardware chip. For example, a full tile connection may be used in one region while a partial tile connection is used in another region. In another example, an N-linear connection arrangement may be used in one region while an N-full tile connection is used in another region.

(89) Alternatively to or in addition with dedicated buses between processor subunits of the memory chip, one or more shared buses may be used to interconnect all (or a subset of) the processor subunits of a distributed processor. Collisions on the shared buses may still be avoided by timing data transfers on the shared buses using code executed by the processor subunits, as explained further below. Additionally with or alternatively to shared buses, configurable buses may be used to dynamically connect processor subunits to form groups of processors units connected to separated buses. For example, the configurable buses may include transistors or other mechanisms that may be controlled by processor subunit to direct data transfers to a selected processor subunit.

(90) In both FIGS. 7A and 7B, the plurality of processor subunits of the processing array is spatially distributed among the plurality of discrete memory banks of the memory array. In other alternative embodiments (not shown), the plurality of processor subunits may be clustered in one or more regions of the substrate, and the plurality of memory banks may be clustered in one or more other regions of the substrate. In some embodiments, a combination of spatial distribution and clustering may be used (not shown). For example, one region of the substrate may include a cluster of processor subunits, another region of the substrate may include a cluster of memory banks, and yet another region of the substrate may include processing arrays distributed amongst memory banks.

(91) One of ordinary skill will recognize that arraying processor groups **600** on a substrate is not an exclusive embodiment. For example, each processor subunit may be associated with at least two dedicated memory banks. Accordingly, processing groups **310a**, **310b**, **310c**, and **310d** of FIG. 3B may be used in lieu of or in combination with processing group **600** to form the processing array and the memory array. Other processing groups including, for example, three, four, or more

dedicated memory banks (not shown) may be used.

(92) Each of the plurality of processor subunits may be configured to execute software code associated with a particular application independently, relative to other processor subunits included in the plurality of processor subunits. For example, as explained below, a plurality of sub-series of instructions may be grouped as machine code and provided to each processor subunit for execution.

(93) In some embodiments, each dedicated memory bank comprises at least one dynamic random access memory (DRAM). Alternatively, the memory banks may comprise a mix of memory types, such as static random access memory (SRAM), DRAM, Flash or the like.

(94) In conventional processors, data sharing between processor subunits is usually performed with shared memory. Shared memory typically requires a large portion of chip area and/or performed a bus that is managed by additional hardware (such as arbiters). The bus results in bottlenecks, as described above. In addition, the shared memory, which may be external to the chip, typically includes cache coherency mechanisms and more complex caches (e.g., L1 cache, L2 cache, and shared DRAM) in order to provide accurate and up-to-date data to the processor subunits. As explained further below, the dedicated buses depicted in FIGS. 7A and 7B allow for hardware chips that are free of hardware management (such as arbiters). Moreover, the use of dedicated memories as depicted in FIGS. 7A and 7B allow for the elimination of complex caching layers and coherency mechanism.

(95) Instead, in order to allow each processor subunit to access data calculated by other processor subunits and/or stored in memory banks dedicated to the other processor subunits, buses are provided whose timing is performed dynamically using code individually executed by each processor subunit. This allows for elimination of most, if not all, bus management hardware as conventionally used. Moreover, complex caching mechanisms are replaced with direct transfers over these buses, resulting in lower latency times during memory reads and writes.

(96) Memory-Based Processing Arrays

(97) As depicted in FIGS. 7A and 7B, a memory chip of the present disclosure may operate independently. Alternatively, memory chips of the present disclosure may be operably connected with one or more additional integrated circuits, such as a memory device (e.g., one or more DRAM banks), a system-on-a-chip, a field-programmable gate array (FPGA), or other processing and/or memory chip. In such embodiments, tasks in a series of instructions executed by the architecture may be divided (e.g., by a compiler, as described below) between processor subunits of the memory chip and any processor subunits of the additional integrated circuit(s). For example, the other integrated circuits may comprise a host (e.g., host **350** of FIG. 3A) that inputs instructions and/or data to the memory chip and receives output therefrom.

(98) In order to interconnect memory chips of the present disclosure with one or more additional integrated circuits, the memory chip may include a memory interface, such as a memory interface complying with a Joint Electron Device Engineering Council (JEDEC) standard or any of its variants. The one or more additional integrated circuits may then connect to the memory interface. Accordingly, if the one or more additional integrated circuits are connected to a plurality of memory chips of the present disclosure, data may be shared between the memory chips through the one or more additional integrated circuits. Additionally or alternatively, the one or more additional integrated circuits may include buses to connect to buses on the memory chips of the present disclosure such that the one or more additional integrated circuits may execute code in tandem with the memory chips of the present disclosure. In such embodiments, the one or more additional integrated circuits further assist with distributed processing even though they may be on different substrates than the memory chips of the present disclosure.

(99) Furthermore, memory chips of the present disclosure may be arrayed in order to form an array of distributed processors. For example, one or more buses may connect a memory chip **770a** to an additional memory chip **770b**, as depicted in FIG. 7C. In the example of FIG. 7C, memory chip **770a** includes processor subunits with one or more corresponding memory banks dedicated to each

processor subunit, e.g.: Processor subunit **730a** is associated with memory bank **720a**, processor subunit **730b** is associated with memory bank **720b**, processor subunit **730e** is associated with memory bank **720c**, and processor subunit **730f** is associated with memory bank **720d**. Buses connect each processor subunit to its corresponding memory bank. Accordingly, bus **740a** connects processor subunit **730a** to memory bank **720a**, bus **740b** connects processor subunit **730b** to memory bank **720b**, bus **740c** connects processor subunit **730e** to memory bank **720c**, and bus **740d** connects processor subunit **730f** to memory bank **720d**. Moreover, bus **750a** connects processor subunit **730a** to processor subunit **750e**, bus **750b** connects processor subunit **730a** to processor subunit **750b**, bus **750c** connects processor subunit **730b** to processor subunit **750f**, and bus **750d** connects processor subunit **730e** to processor subunit **750f**. Other arrangements of memory chip **770a** may be used, for example, as described above.

(100) Similarly, memory chip **770b** includes processor subunits with one or more corresponding memory banks dedicated to each processor subunit, e.g.: Processor subunit **730c** is associated with memory bank **720e**, processor subunit **730d** is associated with memory bank **720f**, processor subunit **730g** is associated with memory bank **720g**, and processor subunit **730h** is associated with memory bank **720h**. Buses connect each processor subunit to its corresponding memory bank. Accordingly, bus **740e** connects processor subunit **730c** to memory bank **720e**, bus **740f** connects processor subunit **730d** to memory bank **720f**, bus **740g** connects processor subunit **730g** to memory bank **720g**, and bus **740h** connects processor subunit **730h** to memory bank **720h**. Moreover, bus **750g** connects processor subunit **730c** to processor subunit **750g**, bus **750h** connects processor subunit **730d** to processor subunit **750h**, bus **750i** connects processor subunit **730c** to processor subunit **750d**, and bus **750j** connects processor subunit **730g** to processor subunit **750h**. Other arrangements of memory chip **770b** may be used, for example, as described above.

(101) The processor subunits of memory chip **770a** and **770b** may be connected using one or more buses. Accordingly, in the example of FIG. 7C, bus **750e** may connect processor subunit **730b** of memory chip **770a** and processor subunit **730c** of memory chip **770b**, and bus **750f** may connect processor subunit **730f** of memory chip **770a** and processor subunit **730c** of memory **770b**. For example, bus **750e** may serve as an input bus to memory chip **770b** (and thus an output bus for memory chip **770a**) while bus **750f** may serve as an input bus to memory chip **770a** (and thus an output bus for memory chip **770b**) or vice versa. Alternatively, buses **750e** and **750f** may both serve as two-way buses between memory chips **770a** and **770b**.

(102) Buses **750e** and **750f** may include direct wires or may be interleaved on a high-speed connection in order to reduce the pins used for the inter-chip interface between memory chip **770a** and integrated circuit **770b**. Moreover, any of the connection arrangements described above used in the memory chip itself may be used to connect the memory chip to one or more additional integrated circuits. For example, memory chip **770a** and **770b** may be connected using a full-tile or partial-tile connection rather than only two buses as shown in FIG. 7C.

(103) Accordingly, although depicted using buses **750e** and **750f**, architecture **760** may include fewer buses or additional buses. For example, a single bus between processor subunits **730b** and **730c** or between processor subunits **730f** and **730c** may be used. Alternatively, additional buses, e.g., between processor subunits **730b** and **730d**, between processor subunits **730f** and **730d**, or the like, may be used.

(104) Furthermore, although depicted as using a single memory chip and an additional integrated circuit, a plurality of memory chips may be connected using buses as explained above. For example, as depicted in the example of FIG. 7C, memory chips **770a**, **770b**, **770c**, and **770d** are connected in an array. Each memory chip includes processor subunits and dedicated memory banks similar to the memory chips described above. Accordingly, a description of these components is not repeated here.

(105) In the example of FIG. 7C, memory chips **770a**, **770b**, **770c**, and **770d** are connected in a loop. Accordingly, bus **750a** connects memory chips **770a** and **770d**, bus **750c** connects memory

chips 770a and 770b, bus 750e connects memory chips 770b and 770c, and bus 750g connects memory chips 770c and 770d. Although memory chips 770a, 770b, 770c, and 770d may be connected with full-tile connections, partial-tile connections, or other connection arrangements, the example of FIG. 7C allows for fewer pin connections between memory chips 770a, 770b, 770c, and 770d.

(106) Relatively Large Memories

(107) Embodiments of the present disclosure may use dedicated memories of relatively large size as compared with shared memories of conventional processors. The use of dedicated memories rather than shared memories allows for gains in efficiency to continue without tapering off with memory increases. This allows for memory-intensive tasks such as neural network processing and database queries to be performed more efficiently than in conventional processors, where the efficiency gains of increasing shared memory taper off due to the von Neumann bottleneck.

(108) For example, in distributed processors of the present disclosure, a memory array disposed on the substrate of the distributed processor may include a plurality of discrete memory banks. Each of the discrete memory banks may have a capacity greater than one megabyte, as well as a processing array disposed on the substrate, including a plurality of processor subunits. As explained above, each one of the processor subunits may be associated with a corresponding, dedicated one of the plurality of discrete memory banks. In some embodiments, the plurality of processor subunits may be spatially distributed among the plurality of discrete memory banks within the memory array. By using dedicated memories of at least one megabyte, rather than shared caches of a few megabytes for a large CPU or GPU, the distributed processors of the present disclosure gain efficiencies that are not possible in conventional systems due to the von Neumann bottleneck in CPUs and GPUs.

(109) Different memories may be used as the dedicated memories. For example, each dedicated memory bank may comprise at least one DRAM bank. Alternatively, each dedicated memory bank may comprise at least one static random access memory bank. In other embodiments, different types of memories may be combined on a single hardware chip.

(110) As explained above, each dedicated memory may be at least one megabyte. Accordingly, each dedicated memory bank may be the same size or at least two of the plurality of memory banks may have different sizes.

(111) Moreover, as described above, the distributed processor may include a first plurality of buses, each connecting one of the plurality of processor subunits to a corresponding, dedicated memory bank and a second plurality of buses, each connecting one of the plurality of processor subunits to another one of the plurality of processor subunits.

(112) Synchronization Using Software

(113) As explained above, hardware chips of the present disclosure may manage data transfers using software rather than hardware. In particular, because the timings of transfers on the buses, reads and writes to the memories, and calculations of the processor subunits are set by the sub-series of instructions executed by the processor subunits, hardware chips of the present disclosure may execute code to prevent collisions on the buses. Accordingly, hardware chips of the present disclosure may avoid hardware mechanisms conventionally used to manage data transfers (such as network controllers within a chip, packet parsers and packets transferors between processor subunits, bus arbitrators, a plurality of buses to avoid arbitration, or the like).

(114) If hardware chips of the present disclosure transferred data conventionally, connecting N processor subunits with buses would require bus arbitration or wide MUXs controlled by an arbiter. Instead, as described above, embodiments of the present disclosure may use a bus that is only a wire, an optical cable, or the like between processor subunits, where the processor subunits individually execute code to avoid collision on the buses. Accordingly, embodiments of the present disclosure may preserve space on the substrate as well as materials cost and efficiency losses (e.g., due to power and time consumption by arbitration). The efficiency and space gains are even greater when compared to other architectures using first-in-first-out (FIFO) controllers and/or mailboxes.

(115) Furthermore, as explained above, each processor subunit may include one or more accelerators in addition to one or more processing elements. In some embodiments, the accelerator(s) may read and write from the buses rather than the processing element(s). In such embodiments, additional efficiency may be obtained by allowing the accelerator(s) to transmit data during the same cycle in which the processing element(s) perform one or more calculations. Such embodiments, however, require additional materials for the accelerator(s). For example, additional transistors may be required for fabrication of the accelerator(s).

(116) The code also may account for the internal behavior, including timing and latencies, of the processor subunits (e.g., including the processing elements and/or accelerators forming part of the processor subunit). For example, a compiler (as described below) may perform pre-processing that accounts for the timing and latencies when generating the sub-series of instructions that control the data transfers.

(117) In one example, a plurality of processor subunits may be assigned a task of calculating a neural network layer containing a plurality of neurons fully-connected to a previous layer of a larger plurality of neurons. Assuming data of the previous layer is evenly spread between the plurality of processor subunits, one way to perform the calculation may be to configure each processor subunit to transmit the data of the previous layer to the main bus in turn and then each processor subunit will multiply this data by the weight of the corresponding neuron that the subunit implements. Because each processor subunit calculates more than one neuron, each processor subunit will transmit the data of the previous layer a number of times equal to the number of neurons. Thus, the code of each processor subunit is not the same as the code for other processor subunits because the subunits will transmit at different times.

(118) In some embodiments, a distributed processor may comprise a substrate (e.g., a semiconductor substrate, such as silicon and/or a circuit board, such as a flexible circuit board) with a memory array disposed on the substrate, the memory array including a plurality of discrete memory banks, and a processing array disposed on the substrate, the processing array including a plurality of processor subunits, as depicted, e.g., in FIGS. 7A and 7B. As explained above, each one of the processor subunits may be associated with a corresponding, dedicated one of the plurality of discrete memory banks. Moreover, as depicted, e.g., in FIGS. 7A and 7B, the distributed processor may further comprise a plurality of buses, each one of the plurality of buses connecting one of the plurality of processor subunits to at least another one of the plurality of processor subunits.

(119) As explained above, the plurality of buses may be controlled in software. Accordingly, the plurality of buses may be free of timing hardware logic components such that data transfers between processor subunits and across corresponding ones of the plurality of buses are uncontrolled by timing hardware logic components. In one example, the plurality of buses may be free of bus arbiters such that data transfers between processor subunits and across corresponding ones of the plurality of buses are uncontrolled by bus arbiters.

(120) In some embodiments, as depicted, e.g., in FIGS. 7A and 7B, the distributed processor may further comprise a second plurality of buses connecting one of the plurality of processor subunits to a corresponding, dedicated memory bank. Similar to the plurality of buses described above, the second plurality of buses may be free of timing hardware logic components such that data transfers between processor subunits and corresponding, dedicated memory banks are uncontrolled by timing hardware logic components. In one example, the second plurality of buses may be free of bus arbiters such that data transfers between processor subunits and corresponding, dedicated memory banks are uncontrolled by bus arbiters.

(121) As used herein, the phrase “free of” does not necessarily imply the absolute absence of components, such as timing hardware logic components (e.g., bus arbiters, arbitration trees, FIFO controllers, mailboxes, or the like). Such components may still be included in a hardware chip described as “free of” those components. Instead, the phrase “free of” refers to the function of the hardware chip; that is, a hardware chip “free of” timing hardware logic components controls the

timing of its data transfers without use of the timing hardware logic components, if any, included therein. For example, a hardware chip that executes code including sub-series of instructions that control data transfers between processor subunits of the hardware chip, even if the hardware chip includes timing hardware logic components as a secondary precaution to protect against collisions due to errors in the executed code.

(122) As explained above, the plurality of buses may comprise at least one of wires or optical fibers between corresponding ones of the plurality of processor subunits. Accordingly, in one example, a distributed processor free of timing hardware logic components may include only wires or optical fibers without bus arbiters, arbitration trees, FIFO controllers, mailboxes, or the like.

(123) In some embodiments, the plurality of processor subunits is configured to transfer data across at least one of the plurality of buses in accordance with code executed by the plurality of processor subunits. Accordingly, as explained below, a compiler may organize sub-series of instructions, each sub-series comprising code executed by a single processor subunit. The sub-series instructions may instruct the processor subunit when to transfer data onto one of the buses and when to retrieve data from the buses. When the sub-series are executed in tandem across the distributed processor, the timing of transfers between the processor subunits may be governed by the instructions to transfer and retrieve included in the sub-series. Thus, the code dictates timing of data transfers across at least one of the plurality of buses. The compiler may generate code to be executed by a single processor subunit. Additionally, the compiler may generate code to be executed by groups of processor subunits. In some cases, the compiler may treat all the processor subunits together as if they were one super-processor (e.g., a distributed processor), and the compiler may generate code for execution by that defined super-processor/distributed processor.

(124) As explained above and depicted in FIGS. 7A and 7B, the plurality of processor subunits may be spatially distributed among the plurality of discrete memory banks within the memory array. Alternatively, the plurality of processor subunits may be clustered in one or more regions of the substrate, and the plurality of memory banks may be clustered in one or more other regions of the substrate. In some embodiments, a combination of spatial distribution and clustering may be used, as explained above.

(125) In some embodiments, a distributed processor may comprise a substrate (e.g., a semiconductor substrate, including silicon and/or a circuit board, such as a flexible circuit board) with a memory array disposed on the substrate, the memory array including a plurality of discrete memory banks. A processing array may also be disposed on the substrate, the processing array including a plurality of processor subunits, as depicted, e.g., in FIGS. 7A and 7B. As explained above, each one of the processor subunits may be associated with a corresponding, dedicated one of the plurality of discrete memory banks. Moreover, as depicted, e.g., in FIGS. 7A and 7B, the distributed processor may further comprise a plurality of buses, each one of the plurality of buses connecting one of the plurality of processor subunits to a corresponding, dedicated one of the plurality of discrete memory banks.

(126) As explained above, the plurality of buses may be controlled in software. Accordingly, the plurality of buses may be free of timing hardware logic components such that data transfers between a processor subunit and a corresponding, dedicated one of the plurality of discrete memory banks and across a corresponding one of the plurality of buses are not controlled by timing hardware logic components. In one example, the plurality of buses may be free of bus arbiters such that data transfers between processor subunits and across corresponding ones of the plurality of buses are uncontrolled by bus arbiters.

(127) In some embodiments, as depicted, e.g., in FIGS. 7A and 7B, the distributed processor may further comprise a second plurality of buses connecting one of the plurality of processor subunits to at least another one of the plurality of processor subunits. Similar to the plurality of buses described above, the second plurality of buses may be free of timing hardware logic components such that data transfers between processor subunits and corresponding, dedicated memory banks are

uncontrolled by timing hardware logic components. In one example, the second plurality of buses may be free of bus arbiters such that data transfers between processor subunits and corresponding, dedicated memory banks are uncontrolled by bus arbiters.

(128) In some embodiments, the distributed processor may use a combination of software timing with hardware timing components. For example, a distributed processor may comprise a substrate (e.g., a semiconductor substrate, including silicon and/or a circuit board, such as a flexible circuit board) with a memory array disposed on the substrate, the memory array including a plurality of discrete memory banks. A processing array may also be disposed on the substrate, the processing array including a plurality of processor subunits, as depicted, e.g., in FIGS. 7A and 7B. As explained above, each one of the processor subunits may be associated with a corresponding, dedicated one of the plurality of discrete memory banks. Moreover, as depicted, e.g., in FIGS. 7A and 7B, the distributed processor may further comprise a plurality of buses, each one of the plurality of buses connecting one of the plurality of processor subunits to at least another one of the plurality of processor subunits. Moreover, as explained above, the plurality of processor subunits may be configured to execute software that controls timing of data transfers across the plurality of buses to avoid colliding data transfers on at least one of the plurality of buses. In such an example, the software may control the timing of the data transfers, but the transfers themselves may be controlled, at least in part, by one or more hardware components.

(129) In such embodiments, the distributed processor may further comprise a second plurality of buses connecting one of the plurality of processor subunits to a corresponding, dedicated memory bank. Similar to the plurality of buses described above, the plurality of processor subunits may be configured to execute software that controls timing of data transfers across the second plurality of buses to avoid colliding data transfers on at least one of the second plurality of buses. In such an example, as explained above, the software may control the timing of the data transfers, but the transfers themselves may be controlled, at least in part, by one or more hardware components.

(130) Division of Code

(131) As explained above, hardware chips of the present disclosure may execute code in parallel across processor subunits included on a substrate forming the hardware chip. Additionally, hardware chips of the present disclosure may perform multitasking. For example, hardware chips of the present disclosure may perform area multitasking, in which one group of processor subunits of the hardware chip execute one task (e.g., audio processing) while another group of processor subunits of the hardware chip execute another task (e.g., image processing). In another example, hardware chips of the present disclosure may perform timing multitasking, in which one or more processor subunits of the hardware chip execute one task during a first period of time and another task during a second period of time. A combination of area and timing multitasking may also be used such that one task may be assigned to a first group of processor subunits during a first period of time while another task may be assigned to a second group of processor subunits during the first period of time, after which a third task may be assigned to processor subunits included in the first group and the second group during a second period of time.

(132) In order to organize machine code for execution on memory chips of the present disclosure, machine code may be divided between processor subunits of the memory chip. For example, a processor on a memory chip may comprise a substrate and a plurality of processor subunits disposed on the substrate. The memory chip may further comprise a corresponding plurality of memory banks disposed on the substrate, each one of the plurality processor subunits being connected to at least one dedicated memory bank not shared by any other processor subunit of the plurality of processor subunits. Each processor subunit on the memory chip may be configured to execute a series of instructions independent from other processor subunits. Each series of instructions may be executed by configuring one or more general processing elements of the processor subunit in accordance with code defining the series of instructions and/or by activating one or more special processing elements (e.g., one or more accelerators) of the processor subunit in

accordance with a sequence provided in the code defining the series of instructions.

(133) Accordingly, each series of instructions may define a series of tasks to be performed by a single processor subunit. A single task may comprise an instruction within an instruction set defined by the architecture of one or more processing elements in the processor subunit. For example, the processor subunit may include particular registers, and a single task may push data onto a register, pull data from a register, perform an arithmetic function on data within a register, perform a logic operation on data within a register, or the like. Moreover, the processor subunit may be configured for any number of operands, such as a 0-operand processor subunit (also called a “stack machine”), a 1-operand processor subunit (also called an accumulator machine), a 2-operand processor subunit (such as a RISC), a 3-operand processor subunit (such as a complex instruction set computer (CISC)), or the like. In another example, the processor subunit may include one or more accelerators, and a single task may activate an accelerator to perform a specific function, such as a MAC function, a MAX function, a MAX-0 function, or the like.

(134) The series of instructions may further include tasks for reading and writing from the dedicated memory banks of the memory chip. For example, a task may include writing a piece of data to a memory bank dedicated to the processor subunit executing the task, reading a piece of data from a memory bank dedicated to the processor subunit executing the task, or the like. In some embodiments, the reading and writing may be performed by the processor subunit in tandem with a controller of the memory bank. For example, the processor subunit may execute a read or write task by sending a control signal to the controller to perform the read or write. In some embodiments, the control signal may include a particular address to use for reads and writes. Alternatively, the processor subunit may defer to the memory controller to select an available address for the reads and writes.

(135) Additionally or alternatively, the reading and writing may be performed by one or more accelerators in tandem with a controller of the memory bank. For example, the accelerators may generate the control signals for the memory controller, similar to how the processor subunit generates control signals, as described above.

(136) In any of the embodiments described above, an address generator may also be used to direct the reads and writes to specific addresses of a memory bank. For example, the address generator may comprise a processing element configured to generate memory addresses for reads and writes. The address generator may be configured to generate addresses in order to increase efficiency, e.g., by writing results of a later calculation to the same address as the results of a former calculation that are no longer needed. Accordingly, the address generator may generate the controls signals for the memory controller, either in response to a command from the processor subunit (e.g., from a processing element included therein or from one or more accelerator(s) therein) or in tandem with the processor subunit. Additionally or alternatively, the address generator may generate the addresses based on some configuration or registers for example generating a nested loop structure to iterate on certain addresses in the memory at a certain pattern.

(137) In some embodiments, each series of instructions may comprise a set of machine code defining a corresponding series of tasks. Accordingly, the series of tasks described above may be encapsulated within machine code comprising the series of instructions. In some embodiments, as explained below with respect to FIG. 8, the series of tasks may be defined by a compiler configured to distribute a higher-level series of tasks amongst the plurality of logic circuits as a plurality of series of tasks. For example, the compiler may generate the plurality of series of tasks based on the higher-level series of tasks such that the processor subunits, executing each corresponding series of tasks in tandem, perform the same function as outlined by the higher-level series of tasks.

(138) As explained further below, the higher-level series of tasks may comprise a set of instructions in a human-readable programming language. Correspondingly, the series of tasks for each processor subunit may comprise lower-level series of tasks, each of which comprises a set of instructions in a machine code.

(139) As explained above with respect to FIGS. 7A and 7B, the memory chip may further comprise a plurality of buses, each bus connecting one of the plurality of processor subunits to at least one other of the plurality of processor subunits. Moreover, as explained above, data transfers on the plurality of buses may be controlled using software. Accordingly, data transfers across at least one of the plurality of buses may be predefined by the series of instructions included in a processor subunit connected to the at least one of the plurality of buses. Therefore, one of the tasks included in the series of instructions may include outputting data to one of the buses or pulling data from one of the buses. Such tasks may be executed by a processing element of the processor subunit or by one or more accelerators included in the processor subunit. In the latter embodiment, the processor subunit may perform a calculation or send a control signal to a corresponding memory bank in the same cycle during which accelerator(s) pull data from or place data on one of the buses.

(140) In one example, the series of instructions included in the processor subunit connected to the at least one of the plurality of buses may include a sending task that comprises a command for the processor subunit connected to the at least one of the plurality of buses to write data to the at least one of the plurality of buses. Additionally or alternatively, the series of instructions included in the processor subunit connected to the at least one of the plurality of buses may include a receiving task that comprises a command for the processor subunit connected to the at least one of the plurality of buses to read data from the at least one of the plurality of buses.

(141) Additionally or alternatively to distribution of code amongst processor subunits, data may be divided between memory banks of the memory chip. For example, as explained above, a distributed processor on a memory chip may comprise a plurality of processor subunits disposed on the memory chip and a plurality of memory banks disposed on the memory chip. Each one of the plurality of memory banks may be configured to store data independent from data stored in other ones of the plurality of memory banks, and each one of the plurality of processor subunits may be connected to at least one dedicated memory bank from among the plurality of memory banks. For example, each processor subunit may have access to one or more memory controllers of one or more corresponding memory banks dedicated to the processor subunit, and no other processor subunit may have access to these corresponding one or more memory controllers. Accordingly, the data stored in each memory bank may be unique to the dedicated processor subunit. Moreover, the data stored in each memory bank may be independent of the memory stored in other memory banks because no memory controllers may be shared between memory banks.

(142) In some embodiments, as described below with respect to FIG. 8, the data stored in each of the plurality of memory banks may be defined by a compiler configured to distribute data amongst the plurality of memory banks. Moreover, the compiler may be configured to distribute data defined in a higher-level series of tasks amongst the plurality of memory banks using a plurality of lower-level tasks distributed amongst corresponding processor subunits.

(143) As explained further below, the higher-level series of tasks may comprise a set of instructions in a human-readable programming language. Correspondingly, the series of tasks for each processor subunit may comprise lower-level series of tasks, each of which comprises a set of instructions in a machine code.

(144) As explained above with respect to FIGS. 7A and 7B, the memory chip may further comprise a plurality of buses, each bus connecting one of the plurality of processor subunits to one or more corresponding, dedicated memory banks from among the plurality of memory banks. Moreover, as explained above, data transfers on the plurality of buses may be controlled using software.

Accordingly, data transfers across a particular one of the plurality of buses may be controlled by a corresponding processor subunit connected to the particular one of the plurality of buses. Therefore, one of the tasks included in the series of instructions may include outputting data to one of the buses or pulling data from one of the buses. As explained above, such tasks may be executed by (i) a processing element of the processor subunit or (ii) one or more accelerators included in the processor subunit. In the latter embodiment, the processor subunit may perform a calculation or use

buses connecting the processor subunit to other processor subunits in the same cycle during which accelerator(s) pull data from or place data on one of the buses connected to the one or more corresponding, dedicated memory banks.

(145) Therefore, in one example, the series of instructions included in the processor subunit connected to the at least one of the plurality of buses may include a sending task. The sending task may comprise a command for the processor subunit connected to the at least one of the plurality of buses to write data to the at least one of the plurality of buses for storage in the one or more corresponding, dedicated memory banks. Additionally or alternatively, the series of instructions included in the processor subunit connected to the at least one of the plurality of buses may include a receiving task. The receiving task may comprise a command for the processor subunit connected to the at least one of the plurality of buses to read data from the at least one of the plurality of buses for storage in the one or more corresponding, dedicated memory banks. Accordingly, the sending and receiving tasks in such embodiments may comprise control signals that are sent, along the at least one of the plurality of buses, to one or more memory controllers of the one or more corresponding, dedicated memory banks. Moreover, the sending and receiving tasks may be executed by one portion of the processing subunit (e.g., by one or more accelerators thereof) concurrently with a calculation or other task executed by another portion of the processing subunit (e.g., by one or more different accelerators thereof). An example of such a concurrent execution may include a MAC-relay command, in which receiving, multiplying, and sending are executed in tandem.

(146) In addition to distributing data amongst the memory banks, particular portions of data may be duplicated across different memory banks. For example, as explained above, a distributed processor on a memory chip may comprise a plurality of processor subunits disposed on the memory chip and a plurality of memory banks disposed on the memory chip. Each one of the plurality of processor subunits may be connected to at least one dedicated memory bank from among the plurality of memory banks, and each memory bank of the plurality of memory banks may be configured to store data independent from data stored in other ones of the plurality of memory banks. Moreover, at least some of the data stored in one particular memory bank from among the plurality of memory banks may comprise a duplicate of data stored in at least another one of the plurality of memory banks. For example, a number, string, or other type of data used in the series of instructions may be stored in a plurality of memory banks dedicated to different processor subunits rather than being transferred from one memory bank to other processor subunits in the memory chip.

(147) In one example, parallel string matching may use data duplication described above. For example, a plurality of strings may be compared to the same string. A conventional processor would compare each string in the plurality to the same string in sequence. On a hardware chip of the present disclosure, the same string may be duplicated across the memory banks such that the processor subunits may compare a separate string in the plurality to the duplicated string in parallel.

(148) In some embodiments, as described below with respect to FIG. 8, the at least some data duplicated across the one particular memory bank from among the plurality of memory banks and the at least another one of the plurality of memory banks is defined by a compiler configured to duplicate data across memory banks. Moreover, the compiler may be configured to duplicate the at least some data using a plurality of lower-level tasks distributed amongst corresponding processor subunits.

(149) Duplication of data may be useful for certain tasks that re-use the same portions of data across different calculations. By duplicating these portions of data, the different calculations may be distributed amongst processor subunits of the memory chip for parallel execution while each processor subunit may store the portions of data in, and access the stored portions from, a dedicated memory bank (rather than pushing and pulling the portions of data across buses connecting the processor subunits). In one example, the at least some data duplicated across the one particular

memory bank from among the plurality of memory banks and the at least another one of the plurality of memory banks may comprise weights of a neural network. In this example, each node in the neural network may be defined by at least one processor subunit from among the plurality of processor subunits. For example, each node may comprise machine code executed by the at least one processor subunit defining the node. In this example, duplication of the weights may allow each processor subunit to execute machine code to effect, at least in part, a corresponding node while only accessing one or more dedicated memory banks (rather than performing data transfers with other processor subunits). Because the timing of reads and writes to the dedicated memory bank(s) are independent of other processor subunits while the timing of data transfers between processor subunits requires timing synchronization (e.g., using software, as explained above), duplication of memory to avoid data transfers between processor subunits may produce further efficiencies in overall execution.

(150) As explained above with respect to FIGS. 7A and 7B, the memory chip may further comprise a plurality of buses, each bus connecting one of the plurality of processor subunits to one or more corresponding, dedicated memory banks from among the plurality of memory banks. Moreover, as explained above, data transfers on the plurality of buses may be controlled using software.

Accordingly, data transfers across a particular one of the plurality of buses may be controlled by a corresponding processor subunit connected to the particular one of the plurality of buses.

Therefore, one of the tasks included in the series of instructions may include outputting data to one of the buses or pulling data from one of the buses. As explained above, such tasks may be executed by (i) a processing element of the processor subunit or (ii) one or more accelerators included in the processor subunit. As further explained above, such tasks may include a sending task and/or a receiving tasks that comprise control signals that are sent, along the at least one of the plurality of buses, to one or more memory controllers of the one or more corresponding, dedicated memory banks.

(151) FIG. 8 depicts a flowchart of a method **800** for compiling a series of instructions for execution on an exemplary memory chip of the present disclosure, e.g., as depicted in FIGS. 7A and 7B. Method **800** may be implemented by any conventional processor, whether generic or special-purpose.

(152) Method **800** may be executed as a portion of a computer program forming a compiler. As used herein, a “compiler” refers to any computer program that converts a higher-level language (e.g., a procedural language, such as C, FORTRAN, BASIC, or the like; an object-oriented language, such as Java, C++, Pascal, Python, or the like; etc.) to a lower-level language (e.g., assembly code, object code, machine code, or the like). The compiler may allow a human to program a series of instructions in a human-readable language, which is then converted to a machine-executable language.

(153) At step **810**, the processor may assign tasks associated with the series of instructions to different ones of the processor subunits. For example, the series of instructions may be divided into subgroups, the subgroups to be executed in parallel across the processor subunits. In one example, a neural network may be divided into its nodes, and one or more nodes may be assigned to separate processor subunits. In this example, each subgroup may comprise a plurality of nodes connected across different layers. Thus, a processor subunit may implement a node from a first layer of the neural network, a node from a second layer connected to the node from the first layer implemented by the same processor subunit, and the like. By assigning nodes based on their connections, data transfers between the processor subunits may be lessened, which may result in greater efficiency, as explained above.

(154) As explained above depicted in FIGS. 7A and 7B, the processor subunits may be spatially distributed among the plurality of memory banks disposed on the memory chip. Accordingly, the assignment of tasks may be, at least in part, a spatial divisional as well as a logical division.

(155) At step **820**, the processor may generate tasks to transfer data between pairs of the processor

subunits of the memory chip, each pair of processor subunits being connected by a bus. For example, as explained above, the data transfers may be controlled using software. Accordingly, processor subunits may be configured to push and pull data on buses at synchronized times. The generated tasks may thus include tasks for performing this synchronized pushing and pulling of data.

(156) As explained above, step **820** may include pre-processing to account for the internal behavior, including timing and latencies, of the processor subunits. For example, the processor may use known times and latencies of the processor subunits (e.g., the time to push data to a bus, the time to pull data from a bus, the latency between a calculation and a push or pull, or the like) to ensure that the generated tasks synchronize. Therefore, the data transfers comprising at least one push by one or more processor subunits and at least one pull by one or more processor subunits may occur simultaneously rather than incurring a delay due to timing differences between the processor subunits, latencies of the processor subunits, or the like.

(157) At step **830**, the processor may group the assigned and generated tasks into the plurality of groups of sub-series instructions. For example, the sub-series instructions may each comprise a series of tasks for execution by a single processor subunit. Therefore, each of the plurality of groups of sub-series instructions may correspond to a different one of the plurality of processor sub-units. Accordingly, steps **810**, **820**, and **830** may result in dividing the series of instructions into a plurality of groups of sub-series instructions. As explained above, step **820** may ensure that any data transfers between the different groups are synchronized.

(158) At step **840**, the processor may generate machine code corresponding to each of the plurality of groups of sub-series instructions. For example, the higher-level code representing sub-series instructions may be converted to lower-level code, such as machine code, executable by corresponding processor subunits.

(159) At step **850**, the processor may assign the generated machine code corresponding to each of the plurality of groups of sub-series instructions to a corresponding one of the plurality of processor subunits in accordance with the division. For example, the processor may label each sub-series instructions with an identifier of the corresponding processor subunit. Thus, when the sub-series instructions are uploaded to a memory chip for execution (e.g., by host **350** of FIG. **3A**), each sub-series may configure a correct processor subunit.

(160) In some embodiments, assigning tasks associated with the series of instructions to the different ones of the processor subunits may depend, at least in part, on a spatial proximity between two or more of the processor subunits on the memory chip. For example, as explained above, efficiency may be increased by lessening the number of data transfers between processor subunits. Accordingly, the processor may minimize data transfers that move data across more than two of processor subunits. Therefore, the processor may use a known layout of the memory chip in combination with one or more optimization algorithms (such as a greedy algorithm) in order to assign sub-series to processor subunits in a way that maximizes (at least locally) adjacent transfers and minimizes (at least locally) transfers to non-neighboring processor subunits.

(161) Method **800** may include further optimizations for the memory chips of the present disclosure. For example, the processor may group data associated with the series of instructions based on the division and assign the data to the memory banks in accordance with the grouping. Accordingly, the memory banks may hold data used for the sub-series instructions assigned to each processor subunit to which each memory bank is dedicated.

(162) In some embodiments, grouping the data may include determining at least a portion of the data to duplicate in two or more of the memory banks. For example, as explained above, some data may be used across more than one sub-series instructions. Such data may be duplicated across the memory banks dedicated to the plurality of processor subunits to which the different sub-series instructions are assigned. This optimization may further reduce data transfers across processor subunits.

(163) The output of method **800** may be input to a memory chip of the present disclosure for execution. For example, a memory chip may comprise a plurality of processor subunits and a corresponding plurality of memory banks, each processor subunit being connected to at least one memory bank dedicated to the processor subunit, and the processor subunits of the memory chip may be configured to execute the machine code generated by method **800**. As explained above with respect to FIG. **3A**, host **350** may input the machine code generated by method **800** to the processor subunits for execution.

(164) Sub-Banks and Sub-Controllers

(165) In conventional memory banks, controllers are provided at the bank level. Each bank includes a plurality of mats, which are typically arranged in a rectangular manner but may be arranged in any geometrical shape. Each mat includes a plurality of memory cells, which are also typically arranged in a rectangular manner but may be arranged in any geometrical shape. Each cell may store a single bit of data (e.g., depending on whether the cell is retained at a high voltage or a low voltage).

(166) An example of this conventional architecture is depicted in FIGS. **9** and **10**. As shown in FIG. **9**, at the bank level, a plurality of mats (e.g., mats **930-1**, **930-2**, **940-1**, and **940-2**) may form bank **900**. In a conventional rectangular organization, bank **900** may be controlled across global wordlines (e.g., wordline **950**) and global bitlines (e.g., bitline **960**). Accordingly, row decoder **910** may select the correct wordline based on an incoming control signal (e.g., a request for a read from an address, a request for a write to an address, or the like) and global sense amplifier **920** (and/or a global column decoder, not shown in FIG. **9**) may select the correct bitline based on the control signal. Amplifier **920** may also amplify any voltage levels from a selected bank during a read operation. Although depicted as using a row decoder for initial selecting and performing amplification along columns, a bank may additionally or alternatively use a column decoder for initial selecting and perform amplification along rows.

(167) FIG. **10** depicts an example of a mat **1000**. For example, mat **1000** may form a portion of a memory bank, such as bank **900** of FIG. **9**. As depicted in FIG. **10**, a plurality of cells (e.g., cells **1030-1**, **1030-2**, and **1030-3**) may form mat **1000**. Each cell may comprise a capacitor, a transistor, or other circuitry that stores at least one bit of data. For example, a cell may comprise a capacitor that is charged to represent a '1' and discharged to represent a '0' or may comprise a flip-flop having a first state representing a '1' and a second state representing a '0.' A conventional mat may comprise, for example, 512 bits by 512 bits. In embodiments where mat **1000** forms a portion of MRAM, ReRAM, or the like, a cell may comprise a transistor, resistor, capacitor or other mechanism for isolating an ion or portion of a material that stores at least one bit of data. For example, a cell may comprise an electrolyte ion, a portion of chalcogenide glass, or the like, having a first state representing a '1' and a second state representing a '0.'

(168) As further depicted in FIG. **10**, in a conventional rectangular organization, mat **1000** may be controlled across local wordlines (e.g., wordline **1040**) and local bitlines (e.g., bitline **1050**). Accordingly, wordline drivers (e.g., wordline driver **1020-1**, **1020-2**, . . . , **1020-x**) may control the selected wordline to perform a read, write, or refresh based on a control signal from a controller associated with the memory bank of which mat **1000** forms a part (e.g., a request for a read from an address, a request for a write to an address, a refresh signal). Moreover, local sense amplifiers (e.g., local amplifiers **1010-1**, **1010-2**, . . . , **1010-x**) and/or local column decoders (not shown in FIG. **10**) may control the selected bitline to perform a read, write, or refresh. The local sense amplifiers may also amplify any voltage levels from a selected cell during a read operation. Although depicted as using a wordline driver for initial selecting and performing amplification along columns, a mat may instead use a bitline driver for initial selecting and perform amplification along rows.

(169) As explained above, a large number of mats are duplicated to form a memory bank. Memory banks may be grouped to form a memory chip. For example, a memory chip may comprise eight to thirty-two memory banks. Accordingly, pairing processor subunits with memory banks on a

conventional memory chip may result in only eight to thirty-two processor subunits. Accordingly, embodiments of the present disclosure may include memory chips with additional sub-bank hierarchy. These memory chips of the present disclosure may then include processor subunits with memory sub-banks used as the dedicated memory banks paired with the processor subunits allowing for a larger number of sub processors, which may then achieve higher parallelism and performance of in-memory computing.

(170) In some embodiments of the present disclosure, the global row decoder and global sense amplifier of bank **900** may be replaced with sub-bank controllers. Accordingly, rather than sending control signals to a global row decoder and a global sense amplifier of the memory bank, a controller of the memory bank may direct the control signal to the appropriate sub-bank controller. The direction may be controlled dynamically or may be hard-wired (e.g., via one or more logic gates). In some embodiments, fuses may be used to indicate the controller of each sub bank or mat whether to block or pass the control signal to the appropriate sub-bank or mat. In such embodiments, faulty sub-banks may thus be deactivated using the fuses.

(171) In one example of such embodiments, a memory chip may include a plurality of memory banks, each memory bank having a bank controller and a plurality of memory sub-banks, each memory sub-bank having a sub-bank row decoder and a sub-bank column decoder for allowing reads and writes to locations on the memory sub-bank. Each sub-bank may comprise a plurality of memory mats, each memory mat having a plurality of memory cells and may have internally local row decoders, column decoders, and/or local sense amplifiers. The sub-bank row decoders and the sub-bank column decoders may process read and write requests from the bank controller or from a sub-bank processor subunit used for in memory computations on the sub-bank memory, as described below. Additionally, each memory sub-bank may further have a controller configured to determine whether to process read requests and write requests from the bank controller and/or to forward them to the next level (e.g., of row and column decoders on a mat) or to block the requests, e.g., to allow an internal processing element or processor subunit to access the memory. In some embodiments, the bank controller may be synchronized to a system clock. However, the sub-bank controllers may be not synchronized to the system clock.

(172) As explained above, the use of sub-banks may allow for the inclusion of a larger number of processor subunits in the memory chip than if processor subunits were paired with memory banks of conventional chips. Accordingly, each sub-bank may further have a processor subunit using the sub-bank as a dedicated memory. As explained above, the processor subunit may comprise a RISC, a CISC, or other general-purpose processing subunit and/or may comprise one or more accelerators. Additionally, the processor subunit may include an address generator, as explained above. In any of the embodiments described above, each processor subunit may be configured to access a sub-bank dedicated to the processor subunit using the row decoder and the column decoder of the sub-bank without using the bank controller. The processor sub-unit associated with the sub-bank may also handle the memory mats (including the decoder and memory redundancy mechanisms, described below) and/or determine whether a read or write request from an upper level (e.g., the bank level or the memory level) is forwarded and handled accordingly.

(173) In some embodiments, the sub-bank controller may further include a register that stores a state of the sub-bank. Accordingly, the sub-bank controller may return an error if the sub-bank controller receives a control signal from the memory controller while the register indicates that the sub-bank is in use. In embodiments where each sub-bank further includes a processor subunit, the register may indicate an error if the processor subunit in the sub-bank is accessing the memory in conflict with an external request from the memory controller.

(174) FIG. **11** shows an example of another embodiment of a memory bank using sub-bank controllers. In the example of FIG. **11**, bank **1100** has a row decoder **1110**, a column decoder **1120**, and a plurality of memory sub-banks (e.g., sub-banks **1170a**, **1170b**, and **1170c**) with sub-bank controllers (e.g., controllers **1130a**, **1130b**, and **1130c**). The sub-bank controllers may include

address resolvers (e.g., resolvers **1140a**, **1140b**, and **1140c**), which may determine whether to pass a request to one or more sub-banks controlled by the sub-bank controller.

(175) The sub-bank controllers may further include one or more logic circuits (e.g., logic **1150a**, **1150b**, and **1150c**). For example, a logic circuit comprising one or more processing elements may allow for one or more operations, such as refreshing of cells in the sub-bank, clearing of cells in the sub-bank, or the like, to be performed without processing requests externally from bank **1100**. Alternatively, the logic circuit may comprise a processor subunit, as explained above, such that the processor sub-unit has any sub-banks controlled by the sub-bank controller as corresponding, dedicated memory. In the example of FIG. **11**, logic **1150a** may have sub-bank **1170a** as a corresponding, dedicated memory, logic **1150b** may have sub-bank **1170b** as a corresponding, dedicated memory, and logic **1150c** may have sub-bank **1170c** as a corresponding, dedicated memory. In any of the embodiments described above, the logic circuits may have buses to the sub-banks. e.g., buses **1131a**, **1131b**, or **1131c**. As further depicted in FIG. **11**, the sub-bank controllers may each include a plurality of decoders, such as a sub-bank row decoder and a sub-bank column decoder for allowing reads and writes, either by a processing element or processor subunit or by a higher-level memory controller issuing commands, to locations on the memory sub-bank(s). For example, sub-bank controller **1130a** includes decoders **1160a**, **1160b**, and **1160c**, sub-bank controller **1130b** includes decoders **1160d**, **1160e**, and **1160f**, and sub-bank controller **1130c** includes decoders **1160g**, **1160h**, and **1160i**. The sub-bank controllers may, based on a request from bank row decoder **1110**, select a wordline using the decoders included in the sub-bank controllers. The described system may allow a processing element or processor subunit of the sub-bank to access the memory without interrupting other banks and even other sub-banks, thereby allowing each sub-bank processor subunit to perform memory computations in parallel with the other sub-bank processor subunits.

(176) Furthermore, each sub-bank may comprise a plurality of memory mats, each memory mat having a plurality of memory cells. For example, sub-bank **1170a** includes mats **1190a-1**, **1190a-2**, . . . , **1190a-x**; sub-bank **1170b** includes mats **1190b-1**, **1190b-2**, . . . , **1190b-x**; and sub-bank **1170c** includes mats **1190c-1**, **1190c-2**, . . . , **1190c-3**. As further depicted in FIG. **11**, each sub-bank may include at least one decoder. For example, sub-bank **1170a** includes decoder **1180a**, sub-bank **1170b** includes decoder **1180b**, and sub-bank **1170c** includes decoder **1180c**. Accordingly, bank column decoder **1120** may select a global bitline (e.g., bitline **1121a** or **1121b**) based on external requests while the sub-bank selected by bank row decoder **1110** may use its column decoder to select a local bitline (e.g., bitline **1181a** or **1181b**) based on local requests from the logic circuit to which the sub-bank is dedicated. Accordingly, each processor subunit may be configured to access a sub-bank dedicated to the processor subunit using the row decoder and the column decoder of the sub-bank without using the bank row decoder and the bank column decoder. Thus, each processor subunit may access a corresponding sub-bank without interrupting other sub-banks. Moreover, sub-bank decoders may reflect accessed data to the bank decoders when the request to the sub-bank is external to the processor subunit. Alternatively, in embodiments where each sub-bank has only one row of memory mats, the local bitlines may be the bitlines of the mat rather than bitlines of the sub-bank.

(177) A combination of embodiments using sub-bank row decoders and sub-bank column decoders with the embodiment depicted in FIG. **11** may be used. For example, the bank row decoder may be eliminated but the bank column decoder retained and local bitlines used.

(178) FIG. **12** shows an example of an embodiment of a memory sub-bank **1200** having a plurality of mats. For example, sub-bank **1200** may represent a portion of sub-bank **1100** of FIG. **11** or may represent an alternative implementation of a memory bank. In the example of FIG. **12**, sub-bank **1200** includes a plurality of mats (e.g., mats **1240a** and **1240b**). Moreover, each mat may include a plurality of cells. For example, mat **1240a** includes cells **1260a-1**, **1260a-2**, . . . , **1260a-x**, and mat **1240b** includes cells **1260b-1**, **1260b-2**, . . . , **1260b-x**.

(179) Each mat may be assigned a range of addresses that will be assigned to the memory cells of the mat. These addresses may be configured at production such that mats may be shuffled around and such that faulted mats may be deactivated and left unused (e.g., using one or more fuses, as explained further below).

(180) Sub-bank **1200** receives read and write requests from memory controller **1210**. Although not depicted in FIG. **12**, requests from memory controller **1210** may be filtered through a controller of sub-bank **1200** and directed to an appropriate mat of sub-bank **1200** for address resolution.

Alternatively, at least a portion (e.g., higher bits) of an address of a request from memory controller **1210** may be transmitted to all mats of sub-bank **1200** (e.g., mats **1240a** and **1240b**) such that each mat may process the full address and the request associated with the address only if the mat's assigned address range includes the address specified in the command. Similar to the sub-bank direction described above, the mat determination may be dynamically controlled or may be hardwired. In some embodiments, fuses may be used to determine the address range for each mat, also allowing for disabling of faulty mats by assigning an illegal address range. Mats may additionally or alternatively be disabled by other common methods or connection of fuses.

(181) In any of the embodiments described above, each mat of the sub-bank may include a row decoder (e.g., row decoder **1230a** or **1230b**) for selection of a wordline in the mat. In some embodiments, each mat may further include fuses and comparators (e.g., **1220a** and **1220b**). As described above, the comparators may allow each mat to determine whether to process an incoming request, and the fuses may allow each mat to deactivate if faulty. Alternatively, row decoders for the bank and/or sub-bank may be used rather than a row decoder in each mat.

(182) Furthermore, in any of the embodiments described above, a column decoder included in the appropriate mat (e.g., column decoder **1250a** or **1250b**) may select a local bitline (e.g., bitline **1251** or **1253**). The local bitline may be connected to a global bitline of the memory bank. In embodiments where the sub-bank has local bitlines of its own, the local bitline of the cell may be further connected to the local bitline of the sub-bank. Accordingly, data in the selected cell may be read through the column decoder (and/or sense amplifier) of the cell, then through the column decoder (and/or sense amplifier) of the sub-bank (in embodiments including a sub-bank column decoder and/or sense amplifier), and then through the column decoder (and/or sense amplifier) of the bank.

(183) Mat **1200** may be duplicated and arrayed to form a memory bank (or a memory sub-bank). For example, a memory chip of the present disclosure may comprise a plurality of memory banks, each memory bank having a plurality of memory sub-banks, and each memory sub-bank having a sub-bank controller for processing reads and writes to locations on the memory sub-bank.

Furthermore, each memory sub-bank may comprise a plurality of memory mats, each memory mat having a plurality of memory cells and having a mat row decoder and a mat column decoder (e.g., as depicted in FIG. **12**). The mat row decoders and the mat column decoders may process read and write requests from the sub-bank controller. For example, the mat decoders may receive all requests and determine (e.g., using a comparator) whether to process the request based on a known address range of each mat, or the mat decoders may only receive requests within the known address range based on selection of a mat by the sub-bank (or bank) controller.

(184) Controller Data Transfers

(185) Any of the memory chips of the present disclosure may also share data using memory controllers (or sub-bank controllers or mat controllers) in addition to sharing data using processing subunits. For example, a memory chip of the present disclosure may comprise a plurality of memory banks (e.g., an SRAM bank, a DRAM bank, or the like), each memory bank having a bank controller, a row decoder, and a column decoder for allowing reads and writes to locations on the memory bank, as well as a plurality of buses connecting each controller of the plurality of bank controllers to at least one other controller of the plurality of bank controllers. The plurality of buses may be similar to the buses connecting the processing subunits, as described above, but connecting

the bank controllers directly rather than through the processing subunits. Furthermore, although described as connecting the bank controllers, buses may additionally or alternatively connect sub-bank controllers and/or mat controllers.

(186) In some embodiments, the plurality of buses may be accessed without interruption of data transfers on main buses of the memory banks connected to one or more processor subunits. Accordingly, a memory bank (or sub-bank) may transmit data to or from a corresponding processor subunit in the same clock cycle as transmitting data to or from a different memory bank (or sub-bank). In embodiments where each controller is connected to a plurality of other controllers, the controllers may be configurable for selection of one other of the other controllers for sending or receiving of data. In some embodiments, each controller may be connected to at least one neighboring controller (e.g., pairs of spatially adjacent controllers may be connected to one another).

(187) Redundant Logic in Memory Circuits

(188) The disclosure is generally directed to a memory chip with primary logic portions for on-chip data processing. The memory chip may include redundant logic portions, which may replace defective primary logic portions to increase the fabrication yield of the chip. Thus, the chip may include on-chip components that allow a configuration of logic blocks in the memory chip based on individual testing of the logic portions. This feature of the chip may increase yields because a memory chip with larger areas dedicated to logic portions is more susceptible to fabrication failures. For example, DRAM memory chips with large redundant logic portions may be susceptible to fabrication issues that reduce yield. However, implementing redundant logic portions may result in increased yield and reliability because it provides a manufacturer or user of DRAM memory chips to turn on or off full logic portions while maintaining the ability of high parallelism. It should be noted that here and throughout the disclosure, example of certain memory types (such as DRAM) may be identified in order to facilitate the explanation of disclosed embodiments. It is to be understood, however, that in such instances the identified memory types are not intended to be limiting. Rather, memory types such as DRAM, Flash, SRAM, ReRAM, PRAM, MRAM, ROM, or any other memory may be used together with the disclosed embodiments even if fewer examples are specifically identified in a certain section of the disclosure.

(189) FIG. 13 is a block diagram of an exemplary memory chip **1300**, consistent with disclosed embodiments. Memory chip **1300** may be implemented as a DRAM memory chip. Memory chip **1300** may also be implemented as any type of memory volatile or non-volatile, such as Flash, SRAM, ReRAM, PRAM, and/or MRAM, etc. Memory chip **1300** may include a substrate **1301** in which an address manager **1302**, a memory array **1304** including a plurality of memory banks, **1304(a,a)** to **1304(z,z)**, a memory logic **1306**, a business logic **1308**, and a redundant business logic **1310** are disposed. Memory logic **1306** and business logic **1308** may constitute primary logic blocks, while redundant business logic **1310** may constitute redundant blocks. In addition, memory chip **1300** may include configuration switches, which may include deactivation switches **1312**, and an activation switches **1314**. Deactivation switches **1312** and activation switches **1314** may also be disposed in the substrate **1301**. In this Application, memory logic **1306**, business logic **1308**, and redundant business logic **1310** may also be collectively referred to as the “logic blocks.”

(190) Address manager **1302** may include row and column decoders or other type of memory auxiliaries. Alternatively, or additionally, address manager **1302** may include a microcontroller or processing unit.

(191) In some embodiments, as shown in FIG. 13, memory chip **1300** may include a single memory array **1304** that may arrange the plurality of memory blocks in a two-dimensional array on substrate **1301**. In other embodiments, however, memory chip **1300** may include multiple memory arrays **1304** and each of the memory arrays **1304** may arrange memory blocks in different configurations. For example, memory blocks in at least one of the memory arrays (also known as memory banks) may be arranged in a radial distribution to facilitate routing between address

manager **1302** or memory logic **1306** to the memory blocks.

(192) Business logic **1308** may be used to do the in-memory computation of an application that is not related to the logic used to manage the memory itself. For example, business logic **1308** may implement functions related to AI such as floating, integer, or MAC operations used as activation functions. In addition, business logic **1308** may implement data base related functions like min, max, sort, count, among others. Memory logic **1306** may perform tasks related to memory management, including (but not limited to) read, write, and refresh operations. Therefore, business logic may be added in one or more of the bank level, mats level, or a group of mats level. Business logic **1308** may have one or more address outputs and one or more data inputs/outputs. For instance, business logic **1308** can address by row/column lines to address manager **1302**. In certain embodiments, however, the logic blocks may be additionally or alternatively addressed via data inputs/outputs.

(193) Redundant business logic **1310** may be a replicate of business logic **1308**. In addition, redundant business logic **1310** may be connected to deactivation switches **1312** and/or activation switches **1314**, which may include small fuse/anti-fuse, and used for logic disabling or enabling one of the instances (e.g., an instance which is connected by default) and enable one of the other logic blocks (e.g., an instance which is disconnected by default). In some embodiments, as further described in connection to FIG. 15, the redundancy of blocks may be local within a logic block, such as business logic **1308**.

(194) In some embodiments, the logic blocks in memory chip **1300** may be connected to subsets of memory array **1304** with dedicated buses. For example, a set of memory logic **1306**, business logic **1308**, and redundant business logic **1310** may be connected to the first row of memory blocks in memory array **1304** (i.e., memory blocks **1304** (a,a) to **1304** (a,z)). The dedicated buses may allow associated logic blocks to quickly access data from the memory blocks without requirements of opening communication lines through, for example, address manager **1302**.

(195) Each of the plurality of primary logic blocks may be connected to at least one of the plurality of memory banks **1304**. Also, redundant blocks, such as redundant business block **1310**, may be connected to at least one of the memory instances **1304**(a,a)-(z,z). Redundant blocks may replicate at least one of the plurality of primary logic blocks, such as memory logic **1306** or business logic **1308**. Deactivation switches **1312** may be connected to at least one of the plurality of primary logic blocks and activation switches **1314** may be connected to at least one of the plurality of redundant blocks.

(196) In these embodiments, upon detecting of a fault associated with one of the plurality of primary logic blocks (memory logic **1306** and/or business logic **1308**), deactivation switches **1312** may be configured to disable the one of the plurality of primary logic blocks. Simultaneously, activation switches **1314** may be configured to enable one of the plurality of redundant blocks, such as redundant logic block **1310**, that replicates the one of the plurality of primary logic blocks.

(197) In addition, activation switches **1314** and deactivation switches **1312**, which may collectively be referred to as “configuration switches,” may include an external input to configure the status of the switch. For instance, activation switches **1314** may be configured so an activation signal in the external input causes a closed switch condition, while deactivation switches **1312** may be configured so a deactivation signal in the external input causes an open switch condition. In some embodiments, all configuration switches in **1300** may be deactivated by default and become activated or enabled after a test indicates an associated logic block is functional and a signal is applied in the external input. Alternatively, in some cases, all configuration switches in **1300** may be enabled by default and may be deactivated or disabled after a test indicates an associated logic block is not functional and a deactivation signal is applied in the external input.

(198) Regardless of whether a configuration switch is initially enabled or disabled, upon detection of a fault associated with an associated logic block, the configuration switch may disable the associated logic block. In cases where the configuration switch is initially enabled, the state of the

configuration switch may be changed to disabled in order to disable the associated logic block. In cases where the configuration switch is initially disabled, the state of the configuration switch may be left in its disabled state in order to disable the associated logic block. For example, the result of an operability test may indicate that a certain logic block is nonoperational or that it fails to operate within certain specifications. In such cases, the logic block may be disabled by not enabling its corresponding configuration switch.

(199) In some embodiments, configuration switches may be connected to two or more logic blocks and may be configured to choose between different logic blocks. For example, a configuration switch may be connected to both business logic **1308** and redundant logic block **1310**.

Configuration switch may enable redundant logic block **1310** while disabling business logic **1308**.

(200) Alternatively, or additionally, at least one of the plurality of primary logic blocks (memory logic **1306** and/or business logic **1308**) may be connected to a subset of the plurality of memory banks or memory instances **1304** with a first dedicated connection. Then, at least one of the plurality of redundant blocks (such as redundant business logic **1310**), which replicates the at least one of the plurality of primary logic blocks, may be connected to the subset of the same plurality of memory banks or instances **1304** with a second dedicated connection.

(201) Moreover, memory logic **1306** may have different functions and capabilities than business logic **1308**. For example, while memory logic **1306** may be designed to enable read and write operations in the memory bank **1304**, business logic **1308** may be designed to perform in-memory computations. Therefore, if the business logic **1308** includes a first business logic block, and the business logic **1308** includes a second business logic block (like redundant business logic **1310**), it is possible to disconnect defective business logic **1308** and reconnect redundant business logic **1310** without missing any capability.

(202) In some embodiments, configuration switches (including deactivation switches **1312** and activation switches **1314**) may be implemented with a fuse, an anti-fuse, or a programmable device (including a one-time programmable device), or other form of non-volatile memory.

(203) FIG. **14** is a block diagram of an exemplary redundant logic block set **1400**, consistent with disclosed embodiments. In some embodiments, redundant logic block set **1400** may be disposed in substrate **1301**. Redundant logic block set **1400** may include at least one of business logic **1308**, and redundant business logic **1310**, connected to switches **1312** and **1314**, respectively. In addition, business logic **1308** and redundant business logic **1310** may be connected to an address bus **1402** and a data bus **1404**.

(204) In some embodiments, as shown in FIG. **14**, the switches **1312** and **1314** may connect logic blocks to a clock node. In this way, the configuration switches may engage or disengage the logic blocks from the clock signal, effectively activating or deactivating the logic blocks. In other embodiments, however, switches **1312** and **1314** may connect logic blocks to other nodes for activation or deactivation. For instance, configuration switches may connect logic blocks to a voltage supply node (e.g., VCC) or to the ground node (e.g., GND) or clock signal. In this way, the logic blocks may be enabled or disabled by the configuration switches because they would create an open circuit or cut-off the logic block power supply.

(205) In some embodiments, as shown in FIG. **14**, address bus **1402** and data bus **1404** may be in opposite sides of the logic blocks, which are connected in parallel to each one of the buses. In this way, routing of the different on-chip components may be facilitated by the logic block set **1400**.

(206) In some embodiments, each one of the plurality of deactivation switches **1312** couple at least one of the plurality of primary logic blocks with a clock node, and each one of the plurality of activation switches **1314** may be couple at least one of the plurality of redundant blocks with the clock node allowing to connect/disconnect the clock as a simple activation/deactivation mechanism.

(207) Redundant business logic **1310** of redundant logic block set **1400** allows the designer to choose, based on area and routing, the blocks that are worth duplication. For example, a chip

designer may select larger blocks for duplication because larger blocks may be more error prone. Thus, a chip designer may decide to duplicate large logic blocks. On the other hand, a designer may prefer to duplicate smaller logic blocks because they are easily duplicated without a significant loss of space. Moreover, using the configuration in FIG. 14, a designer may easily choose to duplicate logic blocks depending on the statistics of errors per area.

(208) FIG. 15 is a block diagram for an exemplary logic block 1500, consistent with disclosed embodiments. The logic block may be business logic 1308 and/or redundant business logic 1310. In other embodiments, however, the exemplary logic block may describe memory logic 1306 or other component of memory chip 1300.

(209) Logic block 1500 presents yet another embodiment where the logic redundancy is used within a small processor pipeline. The logic block 1500 may include a register 1508, a fetch circuit 1504, decoder 1506, and a write-back circuit 1518. In addition, logic block 1500 may include a computation unit 1510 and a duplicated computing unit 1512. However, in other embodiments, logic block 1500 may include other units that do not comprise a controller pipeline but include sporadic processing elements that comprise a required business logic.

(210) Computation unit 1510 and duplicated computation unit 1512 may include a digital circuit capable of performing digital calculations. For example, computation unit 1510 and duplicated computation unit 1512 may include an arithmetic logic unit (ALU) to perform arithmetic and bitwise operations on binary numbers. Alternatively, computation unit 1510 and duplicated computation unit 1512 may include a floating-point unit (FPU), which operates on floating point numbers. In addition, in some embodiments computation unit 1510 and duplicated computation unit 1512 may implement data base related functions like min, max, count, and compare operations, among others.

(211) In some embodiments, as shown in FIG. 15, computation unit 1510 and duplicated computation unit 1512 may be connected to switching circuits 1514 and 1516. When activated the switching circuits may enable or disable the computing units.

(212) In logic block 1500, the duplicated computation unit 1512 may replicate the computation unit 1510. Moreover, in some embodiments, register 1508, fetch circuit 1504, decoder 1506, and write-back circuit 1518 (collectively referred to as the local logic units) may be smaller in size than the computation unit 1510. Because larger elements are more prone to issues during fabrication, a designer may decide to replicate larger units (such as computation unit 1510) instead of smaller units (such as the local logic units). Depending on historic yields and error rates, however, a designer may elect to duplicate local logic units additionally or alternatively to large units (or the entire block). For example, computation unit 1510 may be larger, and thus more error prone, than register 1508, fetch circuit 1504, decoder 1506, and write-back circuit 1518. A designer may choose to duplicate computation unit 1510 instead of the other elements in logic block 1500 or the whole block.

(213) Logic block 1500 may include a plurality of local configuration switches, each one of the plurality of local configuration switches being connected to at least one of the at least one of computation unit 1510 or duplicated computation unit 1512. Local configuration switches may be configured to disable computation unit 1510 and enable duplicated computation unit 1512 when a fault is detected in the computation unit 1510.

(214) FIG. 16 shows block diagrams of exemplary logic blocks connected with a bus, consistent with disclosed embodiments. In some embodiments, logic blocks 1602 (which may represent memory logic 1306, business logic 1308, or redundant business logic 1310) may be independent of each other, may be connected via a bus, and may be activated externally by addressing them specifically. For example, memory chip 1300 may include many logic blocks, each logic block having an ID number. In other embodiments, however, logic blocks 1602 may represent larger units comprised of a plurality one or more of memory logic 1306, business logic 1308, or redundant business logic 1310.

(215) In some embodiments, each one of logic blocks **1602** may be redundant with the other logic blocks **1602**. This complete redundancy, in which all blocks may operate as primary or redundant blocks, may improve fabrication yields because a designer may disconnect faulty units while maintaining functionality of the overall chip. For example, a designer may have the ability to disable logic areas that are prone to errors but maintain similar computation capabilities because the all duplicate blocks may be connected to the same address and data buses. For example, the initial number of logic blocks **1602** may greater than a target capability. Then, disabling some logic blocks **1602** would not affect the target capability.

(216) A bus connected to the logic blocks may include address bus **1614**, command lines **1616**, and data lines **1618**. As shown in FIG. **16**, each one of the logic blocks may be connected independently from each line in the bus. In certain embodiments, however, logic blocks **1602** may be connected in a hierarchical structure to facilitate routing. For instance, each line in the bus may be connected to a multiplexer that routes the line to different logic blocks **1602**.

(217) In some embodiments, to allow external access without knowing the internal chip structure, which may change due to enable and disabled units, each one of the logic blocks may include Fused IDs such as fused identification **1604**. Fused identification **1604** may include an array of switches (like fuses) that determine an ID and may be connected to a managing circuit. For example, fused identification **1604** may be connected to address manager **1302**. Alternatively, fused identification **1604** may be connected to higher memory address units. In these embodiments, fused identification **1604** may be configurable to for a specific address. For example, fused identification **1604** may include a programmable, non-volatile device that determines a final ID based on instructions received form a managing circuit.

(218) A distributed processor on a memory chip may be designed with the configuration depicted in FIG. **16**. A testing procedure executed as BIST at chip wakeup or at factory testing may assign running ID numbers to blocks in the plurality of primary logic blocks (memory logic **1306** and business logic **1308**) that pass a testing protocol. A testing procedure may also assign illegal ID numbers to blocks in the plurality of primary logic blocks that do not pass the testing protocol. The test procedure may also assign running ID numbers to blocks in the plurality of redundant blocks (redundant logic block **1310**) that pass the testing protocol. Because redundant blocks replace failing primary logic blocks, the blocks in the plurality of redundant blocks assigned running ID numbers may be equal to, or greater than, the blocks in the plurality of primary logic blocks assigned illegal ID numbers, thereby disabling the block. In addition, each one of the plurality of primary logic blocks and each one of the plurality of redundant blocks may include at least one fused identification **1604**. Also, as shown in FIG. **16**, the bus connecting logic blocks **1602** may include a command line, a data line, and an address line.

(219) In other embodiments, however, all logic blocks **1602** that are connected to the bus will start disabled and with no ID number. Tested one by one, each good logic block will get a running ID number, and those logic blocks not working will remain with illegal ID, which would disable these blocks. In this manner, redundant logic blocks may improve the fabrication yields by replacing blocks that are known to be defective during the testing process.

(220) Address bus **1614** may couple a managing circuit to each one of the plurality of memory banks, each one of the plurality of primary logic blocks, and each one of the plurality of redundant blocks. These connections allow the managing circuit to, upon detection of the fault associated with a primary logic blocks (such as business logic **1308**), assign an invalid address to the one of the plurality of primary logic blocks and assign a valid address to the one of the plurality of redundant blocks.

(221) For example, as shown in FIG. **16A**, illegal IDs are configured to all logic blocks **1602(a)-(c)** (e.g., address 0xFFF). After testing logic blocks **1602(a)** and **1602(c)** are verified to be functional while logic block **1602(b)** is not functional. In FIG. **16A** unshaded logic blocks may represent logic blocks that passed the functionality test successfully, while shaded logic blocks may represent logic

blocks that failed the test for functionality. Then, the test procedure changes the illegal IDs to legal IDs for logic blocks that are functional while leaving the illegal IDs for logic blocks that are not functional. As an example, in FIG. 16A, the address for logic blocks **1602(a)** and **1602(c)** is changed from 0xFFF to 0x001 and 0x002, respectively. In contrast, the address for logic block **1602(b)** remains the illegal address 0xFFF. In some embodiments, the ID is changed by programming a corresponding fused identification **1604**.

(222) Different results from the testing of logic blocks **1602** may result in a different configuration. For example, as shown in FIG. 16B, address manager **1302** may initially assign illegal IDs to all logic blocks **1602** (i.e., 0xFFF). The testing results, however, may indicate that both logic blocks **1602(a)** and **1602(b)** are functional. In these cases, testing of logic block **1602(c)** may not be necessary because memory chip **1300** may require only two logic blocks. Therefore, to minimize testing resources, logic blocks may be tested only according to the minimum number of functional logic blocks needed by the product definition of **1300**, leaving other logic blocks untested. FIG. 16B also shows unshaded logic blocks, which represent tested logic blocks that passed the test for functionality, and shaded logic blocks, which represent untested logic blocks.

(223) In these embodiments, a production tester (external or internal, automatic or manual) or a controller executing a BIST at startup, may change illegal IDs to running IDs for tested logic blocks that are functional while leaving the illegal IDs to untested logic blocks. As an example, in FIG. 16B, the address for logic blocks **1602(a)** and **1602(b)** is changed from 0xFFF to 0x001 and 0x002, respectively. In contrast, the address for untested logic block **1602(c)** remains with the illegal address 0xFFF.

(224) FIG. 17 is a block diagram for exemplary units **1702** and **1712** connected in series, consistent with disclosed embodiments. FIG. 17 may represent an entire system or chip. Alternatively, FIG. 17 may represent a block in a chip containing other functional blocks.

(225) Units **1702** and **1712** may represent complete units that include a plurality of logic blocks such as memory logic **1306** and/or business logic **1308**. In these embodiments units **1702** and **1712** may also include elements required to perform operations such as address manager **1302**. In other embodiments, however, units **1702** and **1712** may represent logic units such as business logic **1308** or redundant business logic **1310**.

(226) FIG. 17 presents embodiments in which units **1702** and **1712** may need to communicate between themselves. In such cases, units **1702** and **1712** may be connected in series. However, a non-working unit may break the continuity between the logic blocks. Therefore, the connection between units may include a bypass option when a unit needs to be disabled due to a defect. The bypass option can also be a part of the bypassed unit itself.

(227) In FIG. 17 units may be connected in series (e.g., **1702(a)-(c)**), and a failing unit (e.g., **1702(b)**) may be bypassed when it is defective. The units may further be connected in parallel with switching circuits. For example, in some embodiments units **1702** and **1712** may be connected with switching circuits **1722** and **1728**, as depicted in FIG. 17. In the example depicted in FIG. 17, unit **1702(b)** is defective. For example, unit **1702(b)** does not pass a test for a circuit functionality. Therefore, unit **1702(b)** may be disabled using, for example, activation switches **1314** (not shown in FIG. 17) and/or switching circuit **1722(b)** may be activated to bypass unit **1702(b)** and sustain the connectivity between logic blocks.

(228) Accordingly, when a plurality of primary units are connected in series, each one of the plurality of units may be connected in parallel with a parallel switch. Upon detection of a fault associated with the one of the plurality of units, the parallel switch connected to the one of the plurality of units may be activated to connect two of the plurality of units.

(229) In other embodiments, as shown in FIG. 17, switching circuits **1728** may include a sampling point or more that would cause a cycle or cycles delay maintaining synchronization between different lines of units. When a unit is disabled, shorting the connection between adjacent logic blocks may generate synchronization errors with other calculations. For example, if a task requires

data from both A and B lines, and each of A and B is carried by an independent series of units, disabling a unit would cause a desynchronization between the lines that would require further data management. To prevent desynchronizations, sample circuits **1730** may simulate the delay caused by the disabled unit **1712(b)**. Nonetheless, in some embodiments, the parallel switch may include an anti-fuse instead of a sampling circuit **1730**.

(230) FIG. **18** is a block diagram of exemplary units connected in a two-dimension array, consistent with disclosed embodiments. FIG. **18** may represent an entire system or chip. Alternatively, FIG. **18** may represent a block in a chip containing other functional blocks.

(231) Units **1806** may represent autonomous units that include a plurality of logic blocks such as memory logic **1306** and/or business logic **1308**. However, in other embodiments units **1806** may represent logic units such as business logic **1308**. Where convenient, discussion of FIG. **18** may refer to elements identified in FIG. **13** (e.g., memory chip **1300**) and discussed above.

(232) As shown in FIG. **18**, units may be arranged in a two-dimensional array in which units **1806** (which may include or represent one or more of memory logic **1306**, business logic **1308**, or redundant business logic **1310**) are interconnected via switching boxes **1808** and connection boxes **1810**. In addition, in order to control the configuration of the two-dimensional array, the two-dimensional array may include I/O blocks **1804** in the periphery of the two-dimensional array.

(233) Connection boxes **1810** may be programmable and reconfigurable devices that may respond to signals inputted from the I/O blocks **1804**. For example, connection boxes may include a plurality of input pins from units **1806** and may also be connected to switching boxes **1808**. Alternatively, connection boxes **1810** may include a group of switches connecting pins of programmable logic cells with routing tracks, while switching boxes **1808** may include a group of switches connecting different tracks.

(234) In certain embodiments, connection boxes **1810** and switching boxes **1808** may be implemented with configuration switches such as switches **1312** and **1314**. In such embodiments, connection boxes **1810** and switching boxes **1808** may be configured by a production tester or a BIST executed at chip startup.

(235) In some embodiments, connection boxes **1810** and switching boxes **1808** may be configured after units **1806** are tested for a circuit functionality. In such embodiments, I/O blocks **1804** may be used to send testing signals to units **1806**. Depending on the test results, I/O blocks **1804** may send programming signals that configure connection boxes **1810** and switching boxes **1808** in a manner disabling the units **1806** that fail the testing protocol and enabling units **1806** that pass the testing protocol.

(236) In such embodiments, the plurality of primary logic blocks and the plurality of redundant blocks may be disposed on the substrate in a two-dimensional grid. Therefore, each one of the plurality of primary units **1806** and each one of the plurality of redundant blocks, such as redundant business logic **1310**, may be interconnected with switching boxes **1808**, and an input block may be disposed in the periphery of each line and each column of the two-dimensional grid.

(237) FIG. **19** is a block diagram for exemplary units in a complex connection, consistent with disclosed embodiments. FIG. **19** may represent an entire system. Alternatively, FIG. **19** may represent a block in a chip containing other functional blocks.

(238) The complex connection of FIG. **19** includes units **1902(a)-(f)** and configuration switches **1904(a)-(h)**. Units **1902** may represent autonomous units that include a plurality of logic blocks such as memory logic **1306** and/or business logic **1308**. However, in other embodiments units **1902** may represent logic units such as memory logic **1306**, business logic **1308**, or redundant business logic **1310**. Configuration switches **1904** may include any of deactivation switches **1312** and activation switches **1314**.

(239) As shown in FIG. **19**, the complex connection may include units **1902** in two planes. For example, the complex connection may include two independent substrates separated in the z-axis. Alternatively, or additionally, units **1902** may be arranged in two surfaces of a substrate. For

example, with the objective to reduce the area of memory chip **1300**, substrate **1301** may be arranged in two overlapping surfaces and connected with configuration switches **1904** arranged in three dimensions. Configuration switches may include deactivation switches **1312** and/or activation switches **1314**.

(240) A first plane of the substrate may include “main” unit **1902**. These blocks may be enabled by default. In such embodiments, a second plain may include “redundant” unit **1902**. These units may be disabled by default.

(241) In some embodiments, configuration switches **1904** may include anti-fuses. Thus, after testing of units **1902**, the blocks may be connected in a tile of functional units by switching certain anti-fuses to “always-on” and disable selected units **1902**, even if they are in a different plane. In the example presented in FIG. **19**, one of the ‘main’ units (unit **1902(e)**) is not working. FIG. **19** may represent nonfunctional or untested blocks as shaded blocks while tested or functional blocks may be unshaded. Therefore, configuration switches **1904** are configured so one of the logic blocks in a different plane (e.g., unit **1902(f)**) becomes active. In this way even though one of the main logic blocks was defective, the memory chip is still working by replacing a spare logic unit.

(242) FIG. **19** additionally shows that one of the units **1902** (i.e., **1902(c)**) in the second plane is not tested or enabled because the main logic blocks are functional. For example, in FIG. **19**, both main units **1902(a)** and **1902(d)** passed a test for functionality. Thus, units **1902(c)** was not tested or enabled. Therefore, FIG. **19** shows the ability to specifically select the logic blocks that become active depending on testing results.

(243) In some embodiments, as shown in FIG. **19**, not all units **1902** in a first plain may have a corresponding spare or redundant blocks. However, in other embodiments, all units may be redundant with each other for complete redundancy where all units are both primary or redundant. In addition, while some implementations may follow the star network topology depicted in FIG. **19**, other implementation may use parallel connections, serial connections, and/or couple the different elements with configuration switches in parallel or in series.

(244) FIG. **20** is an exemplary flowchart illustrating a redundant block enabling process **2000**, consistent with disclosed embodiments. The enabling process **2000** may be implemented for memory chip **1300** and specially for DRAM memory chips. In some embodiments, process **2000** may include steps of testing each one of a plurality of logic blocks on the substrate of the memory chip for at least one circuit functionality, identifying faulty logic blocks in the plurality of primary logic blocks based on the testing results, testing at least one redundant or additional logic block on the substrate of the memory chip for the at least one circuit functionality, disabling the at least one faulty logic block by applying an external signal to a deactivation switch, and enabling the at least one redundant block by applying the external signal to an activation switch, the activation switch being connected with the at least one redundant block and being disposed on the substrate of the memory chip. The description of FIG. **20** below further elaborates on each step of process **2000**.

(245) Process **2000** may include testing a plurality of logic blocks (step **2002**), such as business block **1308** and a plurality of redundant blocks (e.g., redundant business block **1310**). The testing may be before packaging using, for example, probing stations for on-wafer testing. Step **2000**, however, may also be performed after packaging.

(246) The testing in step **2002** may include applying a finite sequence of testing signals to every logic block in memory chip **1300** or a subset of logic blocks in memory chip **1300**. The testing signals may include requesting a computation that is expected to yield a 0 or a 1. In other embodiments, the testing signal may request reading a specific address in a memory bank or writing in a specific memory bank.

(247) Testing techniques may be implemented to test the response of the logic blocks under iterative processes in step **2002**. For example, the test may involve testing logic blocks by transmitting instructions to write data in a memory bank and then verifying the integrity of the written data. In some embodiments, the testing may include repeating the algorithm with data

inversed.

(248) In alternative embodiments, the testing of step **2002** may include running a model of the logic blocks to generate a target memory image based on a set of testing instructions. Then, the same sequence of instructions may be executed to the logic blocks in the memory chip, and the results may be recorded. The residual memory image of the simulation may also be compared to the image taken from the test, and any mismatch may be flagged as a failure.

(249) Alternatively, in step **2002**, testing may include shadow modeling, where a diagnostic is generated but the results are not necessarily predicted. Instead, the test using shadow modeling may be run in parallel on both the memory chip and a simulation. For example, when the logic blocks in the memory chip complete an instruction or task, the simulation may be signaled to execute the same instruction. Once the logic blocks in the memory chip finalize the instructions, the two models' architectural states may be compared. If there is a mismatch, then a failure is flagged.

(250) In some embodiments, all logic blocks (including, e.g., each one of memory logic **1306**, business logic **1308**, or redundant business logic **1310**) may be tested in step **2002**. In other embodiments, however, only subsets of the logic blocks may be tested in different testing rounds. For example, in a first round of testing only memory logic **1306** and associated blocks may be tested. In a second round, only business logic **1308** and associated blocks may be tested. In a third round, depending on the results of the first two rounds, logic blocks associated with redundant business logic **1310** may be tested.

(251) Process **2000** may continue to step **2004**. In step **2004**, faulty logic blocks may be identified, and faulty redundant blocks may also be identified. For example, logic blocks that do not pass the testing of step **2002** may be identified as faulty blocks in step **2004**. In other embodiments, however, only certain faulty logic blocks may be initially identified. For example, in some embodiments, only logic blocks associated with business logic **1308** may be identified, and faulty redundant blocks are only identified if they are required for substituting a faulty logic block. In addition, identifying faulty blocks may include writing on a memory bank or a nonvolatile memory the identification information of the identified faulty blocks.

(252) In step **2006**, faulty logic blocks may be disabled. For example, using a configuration circuit, the faulty logic blocks may be disabled by disconnecting them from clock, ground, and/or power nodes. Alternatively, faulty logic blocks may be disabled by configuring connection boxes in an arrangement that avoids the logic blocks. Yet, in other embodiments, faulty logic blocks may be disabled by receiving an illegal address from address manager **1302**.

(253) In step **2008**, redundant blocks that duplicate the faulty logic blocks may be identified. To support the same capabilities of the memory chips even though some logic blocks have failed, in step **2008**, redundant blocks that are available and can duplicate faulty logic blocks may be identified. For example, if a logic block that performs multiplications of vectors is determined to be faulty, in step **2008**, an address manager **1302** or an on-chip controller may identify an available redundant logic block that also performs multiplication of vectors.

(254) In step **2010**, the redundant blocks identified in step **2008** may be enabled. In contrast to the disable operation of step **2006**, in step **2010**, the identified redundant blocks may be enabled by connecting them to clock, ground, and/or power nodes. Alternatively, identified redundant blocks may be enabled by configuring connection boxes in an arrangement that connects the identified redundant blocks. Yet, in other embodiments, identified redundant blocks may be enabled by receiving a running address at the test procedure execution time.

(255) FIG. **21** is an exemplary flow chart illustrating an address assignment process **2100**, consistent with disclosed embodiments. The address assignment process **2100** may be implemented for memory chip **1300** and specially for a DRAM memory chips. As described in relation to FIG. **16**, in some embodiments, logic blocks in memory chip **1300** may be connected to a data bus and have an address identification. Process **2100** describes an address assignment method that disables faulty logic blocks and enables logic blocks that pass a test. The steps described in process **2100**

will be described as being performed by a production tester or a BIST executed at chip startup; however, other components of memory chip **1300** and/or external devices may also perform one or more steps of process **2100**.

(256) In step **2102**, the tester may disable all logic and redundant blocks by assigning an illegal identification to each logic block at a chip level.

(257) In step **2104**, the tester may execute a testing protocol of a logic block. For example, the tester may run testing methods described in step **2002** for one or more of the logic blocks in memory chip **1300**.

(258) In step **2106**, depending on the results of the test in step **2104**, the tester may determine whether the logic block is defective. If the logic block is not defective (step **2106**: no), address manager may assign a running ID to the tested logic block in step **2108**. If the logic block is defective (step **2106**: yes), address manager **1302** may leave the illegal ID for the defective logic block in step **2110**.

(259) In step **2112**, address manager **1302** may select a redundant logic block that replicates the defective logic block. In some embodiments, the redundant logic block that replicates the defective logic block may have the same components and connections to the defective logic blocks. In other embodiments, however, the redundant logic block may have different components and/or connections to the defective logic blocks but be able to perform an equivalent operation. For example, if the defective logic block is designed to perform multiplication of vectors, the selected redundant logic block would also be capable of performing multiplication of vectors, even if it does not have the same architecture as the defective unit.

(260) In step **2114**, address manager **1302** may test the redundant block. For instance, the tester may apply the testing techniques applied in step **2104** to the identified redundant block.

(261) In step **2116**, based on the results of testing in step **2114**, the tester may determine whether the redundant block is defective. In step **2118**, if the redundant block is not defective (step **2116**: no), the tester may assign a running ID to the identified redundant block. In some embodiments, process **2100** may return to step **2104** after step **2118**, creating an iteration loop to test all logic blocks in the memory chip.

(262) If the tester determines the redundant block is defective (step **2116**: yes), in step **2120**, the tester may determine if additional redundant blocks are available. For example, the tester may query a memory bank with information regarding available redundant logic blocks. If redundant logic blocks are available (step **2120**: yes), the tester may return to step **2112** and identify a new redundant logic block replicating the defective logic block. If redundant logic blocks are not available (step **2120**: no), in step **2122**, the tester may generate an error signal. The error signal may include information of the defective logic block and the defective redundant block.

(263) Coupled Memory Banks

(264) The presently disclosed embodiments also include a distributed high-performance processor. The processor may include a memory controller that interfaces memory banks and processing units. The processor may be configurable to expedite delivery of data to the processing units for calculations. For example, if a processing unit requires two data instances to perform a task, the memory controller may be configured so communication lines independently provide access to the information from two data instances. The disclosed memory architecture seeks to minimize hardware requirements that are associated with complex cache memory and complex register files schemes. Normally, processor chips include cache hierarchies that allow cores to work directly with registers. However, the cache operations require significant die area and consume additional power. The disclosed memory architecture avoids the use of a cache hierarchy by adding logic components in the memory.

(265) The disclosed architecture also enables strategic (or even optimized) placement of data in memory banks. Even if the memory banks have a single port and high latency, the disclosed memory architecture may enable high performance and avoid memory accessing bottlenecks by

strategically positioning data in different blocks of memory banks. With the goal of providing a continuous stream of data to the processing units, a compilation optimization step may determine how data should be stored in memory banks for specific or generic tasks. Then, the memory controller, which interfaces processing units and memory banks, may be configured to grant access to specific processing units when they require data to perform operations.

(266) The configuration of the memory chip may be performed by a processing unit (e.g., a configuration manager) or an external interface. The configuration may be also written by a compiler or other SW tool. In addition, the configuration of the memory controller may be based on the available ports in the memory banks and the organization of data in the memory banks.

Accordingly, the disclosed architecture may provide processing units with a constant flow of data or simultaneous information from different memory blocks. In this way, computation tasks within the memory may be quickly processed by avoiding latency bottlenecks or cache memory requirements.

(267) Moreover, data stored in the memory chip may be arranged based on compilation optimization steps. The compilation may allow for building of processing routines in which the processor efficiently assigns tasks to processing units without memory latency associated delays. The compilation may be performed by a compiler and transmitted to a host connected to an external interface in the substrate. Normally, high latency for certain access patterns and/or low numbers of ports would result in data bottlenecks for processing units requiring the data. The disclosed compilation, however, may position data in memory banks in a way that enables processing units to continuously receive data even with disadvantageous memory types.

(268) Furthermore, in some embodiments, a configuration manager may signal required processing units based on computations that are required by a task. Different processing units or logic blocks in the chip may have specialized hardware or architectures for different tasks. Therefore, depending on the task that will be performed, a processing unit, or a group of processing units, may be selected to perform the task. The memory controller on the substrate may be configurable to route data, or grant access, according to the selection of processing subunits to improve data transfer rates. For example, based on the compilation optimization and the memory architecture, processing units may be granted access to memory banks when they are required to perform a task.

(269) Moreover, the chip architecture may include on-chip components that facilitate transfer of data by reducing the time required to access data in the memory banks. Therefore, the present disclosure describes chip architecture(s), along with a compilation optimization step, for a high-performance processor capable of performing specific or generic tasks using simple memory instances. The memory instances may have high latency in random access and/or low number of ports, such as those used in a DRAM device or other memory-oriented technologies, but the disclosed architecture may overcome these shortcomings by enabling a continuous (or nearly continuous) flow of data from memory banks to processing units.

(270) In this application, simultaneous communication may refer to communication within a clock cycle. Alternatively, simultaneous communication may refer to sending information within a predetermine amount of time. For example, simultaneous communication may refer to communication within a few nanoseconds.

(271) FIG. 22 provides block diagrams for exemplary processing devices, consistent with disclosed embodiments. FIG. 22A shows a first embodiment of a processing device **2200** in which a memory controller **2210** connects a first memory block **2202** and a second memory block **2204** using multiplexers. Memory controller **2210** may also connect at least a configuration manager **2212**, a logic block **2214**, and multiple accelerators **2216(a)-(n)**. FIG. 22B shows a second embodiment of processing device **2200** in which memory controller **2210** connects memory blocks **2202** and **2204** using a bus that connects memory controller **2210** with at least a configuration manager **2212**, a logic block **2214**, and multiple accelerators **2216(a)-(n)**. In addition, host **2230** may be external and connected to processing device **2200** through, for example, an external interface.

(272) Memory blocks **2202** and **2204** may include a DRAM mats or group of mats, DRAM banks, MRAM\PRAM\RRAM\SRAM units, Flash mats, or other memory technologies. Memory blocks **2202** and **2204** may alternatively include non-volatile memories, a flash memory device, a Resistive Random Access Memory (ReRAM) device, or a Magnetoresistive Random Access Memory (MRAM) device.

(273) Memory blocks **2202** and **2204** may additionally include a plurality of memory cells arranged in rows and columns between a plurality of word lines (not shown) and a plurality of bit lines (not shown). The gates of each row of memory cells may be connected to a respective one of the plurality of word lines. Each column of memory cells may be connected to a respective one of the plurality of bit lines.

(274) In other embodiments, a memory area (including memory blocks **2202** and **2204**) is built from simple memory instances. In this application, the term “memory instance” may be used interchangeably with the term “memory block.” The memory instances (or blocks) may have poor characteristics. For example, the memories may be only one port memories and may have high random-access latency. Alternatively, or additionally, the memories may be inaccessible during column and line changes and face data access problems related to, for example, capacity charging and/or circuitry setups. Nonetheless, the architecture presented in FIG. **22** still facilitates parallel processing in the memory device by allowing dedicated connections between memory instances and processing units and arranging the data in a certain manner that takes the characteristics of the blocks into account.

(275) In some device architectures, memory instances may include several ports, facilitating the parallel operations. Nonetheless, in such embodiments, the chip may still achieve an improved performance when data is compiled and organized based on the chip architecture. For example, a compiler may improve the efficiency of access in the memory area by providing instructions and organizing data placement, so it can be readily access even using one-port memories.

(276) Furthermore, memory blocks **2202** and **2204** may be multiple types for memory in a single chip. For example, memory blocks **2202** and **2204** may be eFlash and eDRAM. Also, memory blocks may include DRAM with instances of ROM.

(277) Memory controller **2210** may include a logic circuit to handle the memory access and return the results to the rest of the modules. For example, memory controller **2210** may include an address manager and selection devices, such as multiplexers, to route data between the memory blocks and processing units or grant access to the memory blocks. Alternatively, Memory controller **2210** may include double data rate (DDR) memory controllers used to drive DDR SDRAM, where data is transferred on both rising and falling edges of the system's memory clock.

(278) In addition, memory controller **2210** may constitute Dual Channel memory controllers. The incorporation of dual channel memory may facilitate control of parallel access lines by memory controller **2210**. The parallel access lines may be configured to have identical lengths to facilitate synchronization of data when multiple lines are used in conjunction. Alternatively, or additionally, the parallel access lines may allow access of multiple memory ports of the memory banks.

(279) In some embodiments processing device **2200** may include one or more muxes that may be connected to processing units. The processing units may include configuration manager **2212**, logic block **2214**, and accelerators **2216**, which may be connected directly to the mux. Also, memory controller **2210** may include at least one data input from a plurality of memory banks or blocks **2202** and **2204** and at least one data output connected to each one of the plurality of processing units. With this configuration, memory controller **2210** may simultaneously receive data from memory banks or memory blocks **2202** and **2204** via the two data inputs, and simultaneously transmits data received via to the at least one selected processing unit via the two data outputs. In some embodiments, however, the at least one data input and at least one data output may be implemented in a single port allowing only read or write operations. In such embodiments, the single port may be implemented as a data bus including data, address, and command lines.

(280) Memory controller **2210** may be connected to each one of the plurality of memory blocks **2202** and **2204**, and may also connect to processing units via, for example, a selection switch. Also processing units on the substrate, including configuration manager **2212**, logic block **2214**, and accelerators **2216**, may be independently connected to memory controller **2210**. In some embodiments, configuration manager **2212** may receive an indication of a task to be performed and, in response, configure memory controller **2210**, accelerators **2216**, and/or logic blocks **2214** according to a configuration stored in memory or supplied externally. Alternatively, memory controller **2210** may be configured by an external interface. The task may require at least one computation that may be used to select at least one selected processing unit from the plurality of processing units. Alternatively, or additionally, the selection may be based at least in part upon a capability of the selected processing unit for performing the at least one computation. In response, memory controller **2210** may grant access to the memory banks, or route data between the at least one selected processing unit and at least two memory banks, using dedicated buses and/or in a pipelined memory access.

(281) In some embodiments, first memory block **2202** of at least two memory blocks may be arranged on a first side of the plurality of processing units; and second memory bank **2204** of the at least two memory banks may be arranged on a second side of the plurality of processing units opposite to the first side. Further, a selected processing unit to perform the task, for instance accelerator **2216(n)**, may be configured to access second memory bank **2204** during a clock cycle in which a communication line is opened to the first memory bank or first memory block **2202**. Alternatively, the selected processing unit may be configured to transfer data to second memory block **2204** during a clock cycle in which a communication line is opened to first memory block **2202**.

(282) In some embodiments, memory controller **2210** may be implemented as an independent element, as shown in FIG. **22**. In other embodiments, however, memory controller **2210** may be embedded in the memory area or may be disposed along accelerators **2216(a)-(n)**.

(283) A processing area in processing device **2200** may include configuration manager **2212**, logic block **2214**, and accelerators **2216(a)-(n)**. Accelerators **2216** may include multiple processing circuits with pre-defined functions and may be defined by a specific application. For example, an accelerator may be a vector multiply accumulate (MAC) unit or a Direct Memory Access (DMA) unit handling memory moving between modules. Accelerators **2216** may also be able to calculate their own address and request the data from memory controller **2210** or write data to it. For example, configuration manager **2212** may signal at least one of accelerators **2216** that he can access the memory bank. Then accelerators **2216** may configure memory controller **2210** to route data or grant access to themselves. In addition, accelerators **2216** may include at least one arithmetic logic unit, at least one vector handling logic unit, at least one string compare logic unit, at least one register, and at least one direct memory access.

(284) Configuration manager **2212** may include digital processing circuits to configure accelerators **2216** and instructs execution of tasks. For example, configuration manager **2212** may be connected to memory controller **2210** and each one of the plurality of accelerators **2216**. Configuration manager **2212** may have its own dedicated memory to hold the configurations of accelerators **2216**. Configuration manager **2212** may use the memory banks to fetch commands and configurations via memory controller **2210**. Alternatively, configuration manager **2212** may be programmed through an external interface. In certain embodiments, configuration manager **2212** may be implemented with an on-chip reduced instruction set computer (RISC) or an on-chip complex CPU with its own cache hierarchy. In some embodiments, configuration manager **2212** may also be omitted and the accelerators can be configured through an external interface.

(285) Processing device **2200** may also include an external interface (not shown). The external interface allows access to the memory from an upper level, such a memory bank controller which receives the command from external host **2230** or on-chip main processor or access to the memory

from external host **2230** or on-chip main processor. The external interface may allow programming of the configuration manager **2212** and the accelerators **2216** by writing configurations or code to the memory via memory controller **2210** to be used later by configuration manager **2212** or the units **2214** and **2216** themselves. The external interface, however, may also directly program processing units without being routed through memory controller **2210**. In case configuration manager **2212** is a microcontroller, configuration manager **2212** may allow loading of code from a main memory to the controller local memory via the external interface. Memory controller **2210** may be configured to interrupt the task in response to receiving a request from the external interface.

(286) The external interface may include multiple connectors associated with logic circuits that provide a glue-less interface to a variety of elements on the processing device. The external interface may include: Data I/O Inputs for data reads and output for data writes; External address outputs; External CEO chip select pins; Active-low chip selectors; Byte enable pins; a pin for wait states on the memory cycle; a Write enable pin; an Output enable-active pin; and read-write enable pin. Therefore, the external interface has the required inputs and outputs to control processes and obtain information from the processing device. For example, the external interface may conform to JEDEC DDR standards. Alternatively, or additionally, external interface may conform to other standards such as SPRO\OSPI or UART.

(287) In some embodiments, the external interface may be disposed on the chip substrate and may be connected external host **2230**. The external host may gain access to memory blocks **2202** and **2204**, memory controller **2210**, and processing units via the external interface. Alternatively, or additionally, external host **2230** may read and write to the memory or may signal configuration manager **2212**, through read and write commands, to perform operations such as starting a process and/or stopping a process. In addition, external host **2230** may configure the accelerators **2216** directly. In some embodiments, external host **2230** be able to perform read/write operations directly on memory blocks **2202** and **2204**.

(288) In some embodiments, configuration manager **2212** and accelerators **2216** may be configured to connect the device area with the memory area using direct buses depending on the target task. For example, a subset of accelerators **2216** may connect with memory instances **2204** when the subset of accelerators has the capability to perform computations required to execute the task. By doing such a separation, it is possible to assure that dedicated accelerators get the bandwidth (BW) needed to memory blocks **2202** and **2204**. Moreover, this configuration with dedicated buses may allow splitting a large memory to smaller instances or blocks because connecting memory instances to memory controller **2210** allows quick access to data in different memories even with high row latency time. To achieve the parallelization of connection, memory controller **2210** may be connected to each of the memory instances with data, address, and/or control buses.

(289) The above-discussed inclusion of memory controller **2210** may eliminate the requirement of a cache hierarchy or complex register tile in the processing device. Although the cache hierarchy can be added to give added capabilities, the architecture in processing device processing device **2200** may allow a designer to add enough memory blocks or instances based on the processing operations and manage the instances accordingly without a cache hierarchy. For example, the architecture in processing device processing device **2200** may eliminate requirements of a cache hierarchy by implementing a pipelined memory access. In the pipelined memory access, processing units may receive a sustaining flow of data in every cycle certain data lines may be opened (or activated) while other data lines receive or transmit data. The sustained flow of data using independent communication lines may allow an improved execution speed and minimum latency due to line changes.

(290) Moreover, the disclosed architecture in FIG. 22 enables a pipelined memory access it may be possible to organize data in a low number of memory blocks and save power losses caused by line switching. For example, in some embodiments, a compiler may communicate host **2230** the

organization of, or a method to organize, data in memory banks to facilitate access to data during a given task. Then, configuration manager **2212** may define which memory banks, and in some cases which ports of the memory banks, may be accessed by the accelerators. This synchronization between the location of data in memory banks and the access method to data, improves computing tasks by feeding data to the accelerators with minimum latency. For example, in embodiments in which configuration manager **2212** includes a RISC/CPU, the method may be implemented in offline software (SW) and then the configuration manager **2212** may be programmed to execute the method. The method may be developed in any language executable by RISC/CPU computers and may be executed on any platform. The inputs of the method may include configuration of the memories behind memory controller and the data itself along with the pattern of memory accesses. In addition, the method may be implemented in a language or machine language specific to the embodiment and may also be just a series of configuration values in binary or text.

(291) As discussed above, in some embodiments, a compiler may provide instructions to host **2230** for organizing data in memory blocks **2202** and **2204** in preparation of a pipelined memory access. The pipelined memory access may generally include steps of: receiving a plurality of addresses of a plurality of memory banks or memory blocks **2202** and **2204**; accessing the plurality of memory banks according to the received addresses using independent data lines; supplying data from a first address through a first communication line to at least one of the plurality of processing units and opening a second communication line to a second address, the first address being in a first memory bank of the plurality of memory banks, the second address being in second memory bank **2204** of the plurality of memory banks; and supplying data from the second address through the second communication line to the at least one of the plurality of processing units and opening a third communication line to a third address in the first memory bank in the first line within a second clock cycle. In some embodiments, the pipelined memory access may be executed with two memory blocks being connected to a single port. In such embodiments, memory controller **2210** may hide the two memory blocks behind a single port but transmit data to the processing units with the pipelined memory access approach.

(292) In some embodiments, a compiler can run on host **2230** before executing a task. In such embodiments, the compiler may be able to determine a configuration of data flow based on the architecture of the memory device since the configuration would be known to the compiler.

(293) In other embodiments, if the configuration of memory blocks **2204** and **2202** is unknown at offline time, the pipelined method can run on host **2230** which may arrange data in memory blocks before starting calculations. For example, host **2230** may directly write data in memory blocks **2204** and **2202**. In such embodiments, processing units, such as configuration manager **2212** and memory controller **2210** may not have information regarding required hardware until run time. Then, it may be necessary to delay the selection of an accelerator **2216** until a task starts running. In these situations, the processing units or memory controller **2210** may randomly select an accelerator **2216** and create a test data access pattern, which may be modified as the task is executed.

(294) Nonetheless, when the task is known in advance, a compiler may organize data and instructions in memory banks for host **2230** to provide to a processing unit, such as configuration manager **2212**, to set signal connections that minimize access latency. For example, in some cases n words may be needed at the same time by accelerators **2216**. However, each memory instance supports retrieving only m words at a time, where “ m ” and “ n ” are integers and $m < n$. Thus, the compiler may place the needed data across different memory instances or blocks facilitating data access. Also, to avoid line miss latencies, a host may split data in different lines of different memory instances if processing device **2200** includes multiple memory instances. The division of data may allow accessing the next line of data in the next instance while still using data from the current instance.

(295) For example, accelerator **2216(a)** may be configured to multiply two vectors. Each one of the

vectors may be stored in independent memory blocks, such as memory blocks **2202** and **2204**, and each vector may include multiple words. Therefore, to complete a task requiring a multiplication by accelerator **2216(a)**, it may be necessary to access the two memory blocks and retrieve multiple words. However, in some embodiments, memory blocks only allow access of one word per clock cycle. For instance, memory blocks may have a single port. In these cases, to expedite data transmittal during an operation, a compiler may organize the words composing vectors in different memory blocks allowing parallel and/or simultaneous reading of the words. In these situations, a compiler may store words in memory blocks that have a dedicated line. For instance, if each vector includes two words and memory controller has direct access to four memory blocks, a compiler may arrange data in four memory blocks, each one transmitting a word and expediting data delivery. Moreover, in embodiments when memory controller **2210** may have more than a single connection to each memory block, the compiler may instruct configuration manager **2212** (or other processing unit) to access ports specific ports. In this way, processing device **2200** may perform a pipelined memory access, continuously providing data to processing units by simultaneously loading words in some lines and transmitting data in other lines. Thus, this pipelined memory access avoid may avoid latency issues.

(296) FIG. **23** is a block diagram of an exemplary processing device **2300**, consistent with disclosed embodiments. The block diagram shows a simplified processing device **2300** displaying a single accelerator in the form of MAC Unit **2302**, configuration manager **2304** (equivalent or similar to configuration manager **2212**), memory controller **2306** (equivalent or similar to memory controller **2210**), and a plurality of memory blocks **2308(a)-(d)**.

(297) In some embodiments, MAC unit **2302** may be a specific accelerator for processing a particular task. By way of example, the processing device **2300** may be tasked with 2D-convolutions. Then, configuration manager **2304** can signal an accelerator that has the appropriate hardware to perform calculations associated with the task. For instance, MAC unit **2302** may have four internal incrementing counters (logical adders and registers to manage the four loops needed by a convolution calculation) and a multiply accumulate unit. Configuration manager **2304** may signal MAC unit **2302** to process incoming data and execute the task. Configuration manager **2304** may transmit an indication to MAC unit **2302** to execute the task. In these situations, MAC unit **2302** may iterate over calculated addresses, multiply the numbers, and accumulate them to an internal register.

(298) In some embodiments, configuration manager **2304** may configure the accelerators while memory controller **2306** grants access to blocks **2308** and MAC unit **2302** using dedicated buses. In other embodiments, however, memory controller **2306** can directly configure the accelerators based on instructions received from configuration manager **2304** or an external interface. Alternatively, or additionally, configuration manager **2304** can pre-load a few configurations and allow the accelerator to iteratively run on different addresses with different sizes. In such embodiments, configuration manager **2304** may include a cache memory that stores a command before it is transmitted to at least one of the plurality of processing units, such as accelerators **2216**. However, in other embodiments configuration manager **2304** may not include a cache.

(299) In some embodiments, configuration manager **2304** or memory controller **2306** may receive addresses that need to be accessed for a task. Configuration manager **2304** or memory controller **2306** may check a register to determine whether the address is already in a loaded line to one of memory blocks **2308**. If so, memory controller **2306** may read the word from memory block **2308** and pass it to the MAC unit **2302**. If the address is not in a loaded line, configuration manager **2304** may request memory controller **2306** may load the line and signal MAC unit **2302** to delay until it is retrieved.

(300) In some embodiments, as shown in FIG. **23**, memory controller **2306** may include two inputs form two independent addresses. But if more than two addresses should be accessed simultaneously, and these addresses are in a single memory block (for example it is only in of

memory blocks **2308(a)**), memory controller **2306** or configuration manager **2304** may raise an exception. Alternatively, configuration manager **2304** may return invalid data signal when the two addresses can only be access through a single line. In other embodiments, the unit may delay the process execution until it is possible to retrieve all needed data. This may diminish the overall performance. Nonetheless, a compiler may be able to find a configuration and data placement that would prevent delays.

(301) In some embodiments, a compiler may create a configuration or instruction set for processing device **2300** that may configure configuration manager **2304** and memory controller **2306** and accelerator **2302** to handle situations in which multiple addresses need to be accessed from a single memory block but the memory block has one port. For instance, a compiler may re-arrange data in memory blocks **2308** such that processing units may access multiple lines in memory blocks **2308**.

(302) In addition, memory controller **2306** may also work simultaneously on more than one input at the same time. For example, memory controller **2306** may allow accessing one of memory blocks **2308** through one port and supplying the data while receiving a request from a different memory block in another input. Therefore, this operation may result in an accelerator **2216** tasked with the exemplary 2D-convolutions receiving data from dedicated lines of communication with the pertinent memory blocks.

(303) Additionally, or alternatively, memory controller **2306** or a logic block may hold refresh counters for every memory block **2308** and handle the refresh of all lines. Having such a counter allows memory controller **2306** to slip in the refresh cycles between dead access times from the devices.

(304) Furthermore, memory controller **2306** may be configurable to perform the pipelined memory access, receiving addresses and opening lines in memory blocks before supplying the data. The pipelined memory access may provide data to processing units without interruption or delayed clock cycles. For example, while memory controller **2306** or one of the logic blocks access data with the right line in FIG. 23, it may be transmitting data in the left line. These methods will be explained in greater detail in connection to FIG. 26.

(305) In response to the required data, processing device **2300** may use multiplexors and/or other switching devices to choose which device gets serviced to perform a given task. For example, configuration manager **2304** may configure multiplexers so at least two data lines reach the MAC unit **2302**. In this way, a task requiring data from multiple addresses, such as 2D-convolutions, may be performed faster because the vectors or words requiring multiplication during convolution can reach the processing unit simultaneously, in a single clock. This data transferring method may allow the processing units, such as accelerators **2216**, to quickly output a result.

(306) In some embodiments, configuration manager **2304** may be configurable to execute processes based on priority of tasks. For example, configuration manager **2304** can be configured to let a running process finish without any interruptions. In that case, configuration manager **2304** may provide an instruction or configurations of a task to accelerators **2216**, let them run uninterrupted, and switch multiplexers only when the task is finished. However, in other embodiments, configuration manager **2304** may interrupt a task and reconfigure data routing when it receives a priority task, such a request from an external interface. Nevertheless, with enough memory blocks **2308**, memory controller **2306** may be configurable to route data, or grant access, to processing units with dedicated lines that do not have to be changed until a task is completed. Moreover, in some embodiments, all devices may be connected by buses to the entries of configuration manager **2304**, and the devices may manage access between themselves and the buses (e.g., using the same logic as a multiplexer). Therefore, memory controller **2306** may be directly connected to a number of memory instances or memory blocks.

(307) Alternatively, memory controller **2306** may be connected directly to memory sub-instances. In some embodiments, each memory instance or block can be built from sub-instances (for example, DRAM may be built from mats with independent data lines arranged in multiple sub-

blocks). Further, the instances may include at least one of DRAM mats, DRAM, banks, flash mats, or SRAM mats or any other type of memory. Then, memory controller **2306** may include dedicated lines to address sub-instances directly to minimize latency during a pipelined memory access. (308) In some embodiments, memory controller **2306** may also hold the logic needed for a specific memory instance (such as row\col decoders, refresh logic, etc.) and memory blocks **2308** may handle its own logic. Therefore, memory blocks **2308** may get an address and generate commands for return\write data.

(309) FIG. **24** depicts exemplary memory configuration diagrams, consistent with disclosed embodiments. In some embodiments, a compiler generating code or configuration for processing device **2200** may perform a method to configure loading from memory blocks **2202** and **2204** by pre-arranging data in each block. For example, a compiler may prearrange data so each word required for a task is correlated to a line of memory instance or memory block(s). But for tasks that require more memory blocks than the one available in processing device **2200**, a compiler may implement methods of fitting data in more than one memory location of each memory block. The compiler may also store data in sequence and evaluate the latency of each memory block to avoid line miss latency. In some embodiments, the host may be part of a processing unit, such as configuration manager **2212**, but in other embodiments the compiler host may be connected to processing device **2200** via an external interface. In such embodiments, the host may run compiling functions, such as the ones described for the compiler.

(310) In some embodiments, configuration manager **2212** may be a CPU or a micro controller (uC). In such embodiments, configuration manager **2212** may have to access the memory to fetch commands or instructions placed in the memory. A specific compiler may generate the code and place it in the memory in a manner that allows for consecutive commands to be stored in the same memory line and across a number of memory banks to allow for the pipelined memory access also on the fetched command. In these embodiments, configuration manager **2212** and memory controller **2210** may be capable of avoiding row latency in linear execution by facilitating the pipelined memory access.

(311) The previous case of linear execution of a program described a method for a compiler to recognize and place the instructions to allow for pipelined memory execution. However other software structures may be more complex and would require the compiler to recognize them and act accordingly. For example, in case a task requires loops and branches, a compiler may place all the loop code inside a single line so that the single line can be looped without line opening latency. Then, memory controller **2210** may not need to change lines during an execution.

(312) In some embodiments, configuration manager **2212** may include internal caching or small memory. The internal caching may store commands that are executed by configuration manager **2212** to handle branches and loops. For example, commands in internal caching memory may include instructions to configure accelerators for accessing memory blocks.

(313) FIG. **25** is an exemplary flowchart illustrating a possible memory configuration process **2500**, consistent with disclosed embodiments. Where convenient in describing memory configuration process **2500**, reference may be made to the identifiers of elements depicted in FIG. **22** and described above. In some embodiments, process **2500** may be executed by a compiler that provides instructions to a host connected through an external interface. In other embodiments, process **2500** may be executed by components of processing device **2200**, such as configuration manager **2212**.

(314) In general, process **2500** may include determining a number of words required simultaneously to perform the task; determining a number of words that can be accessed simultaneously from each one of the plurality of memory banks; and dividing the number of words required simultaneously between multiple memory banks when the number of words required simultaneously is greater than the number of words that can be accessed simultaneously. Moreover, dividing the number of words required simultaneously may include executing a cyclic organization

of words and sequentially assigning one word per memory bank.

(315) More specifically, process **2500** may begin with step **2502**, in which a compiler may receive a task specification. The specification includes required computations and/or a priority level.

(316) In step **2504**, a compiler may identify an accelerator, or group of accelerators, that may perform the task. Alternatively, the compiler may generate instructions so the processing units, such as configuration manager **2212**, may identify an accelerator to perform the task. For example, using the required computation configuration manager **2212** may identify accelerators in the group of accelerators **2216** that may process the task.

(317) In step **2506**, the compiler may determine a number of words that needs to be simultaneously accessed to execute the task. For example, the multiplication of two vectors requires access to at least two vectors, and the compiler may therefore determine that vector words must be simultaneously accessed to perform the operation.

(318) In step **2508**, the compiler may determine a number of cycles necessary to execute the task. For example, if the task requires a convolution operation of four by-products, the compiler may determine that at least 4 cycles will be necessary to perform the task.

(319) In step **2510**, the compiler may place words that are needed to be accessed simultaneously in different memory banks. In that way, memory controller **2210** may be configured to open lines to different memory instances and access the required memory blocks within a clock cycle, without any required cached data.

(320) In step **2512**, the compiler place words that are accessed sequentially in the same memory banks. For example, in the case that four cycles of operations are required, the compiler may generate instructions to write needed words in sequential cycles in a single memory block to avoid changing lines between different memory blocks during execution.

(321) In step **2514**, compiler generate instructions for programing processing units, such as configuration manager **2212**. The instructions may specify conditions to operate a switching device (such as a multiplexor) or configure a data bus. With such instructions, configuration manager **2212** may configure memory controller **2210** to route data from, or grant access to, memory blocks to processing units using dedicated lines of communication according to a task.

(322) FIG. **26** is an exemplary flowchart illustrating a memory read process **2600**, consistent with disclosed embodiments. Where convenient in describing memory read process **2600**, reference may be made to the identifiers of elements depicted in FIG. **22** and described above. In some embodiments, as described below, process **2600** may be implemented by memory controller **2210**. In other embodiments, however, process **2600** may be implemented by other elements in the processing device **2200**, such as configuration manager **2212**.

(323) In step **2602**, memory controller **2210**, configuration manager **2212**, or other processing units may receive an indication to route data from, or grant access to, a memory bank. The request may specify an address and a memory block.

(324) In some embodiments, the request may be received via a data bus specifying a read command in line **2218** and address in line **2220**. In other embodiments, the request may be received via demultiplexers connected to memory controller **2210**.

(325) In step **2604**, configuration manager **2212**, a host, or other processing units, may query an internal register. The internal register may include information regarding opened lines to memory banks, opened addresses, opened memory blocks, and/or upcoming tasks. Based on the information in the internal register, it may be determined whether there are lines opened to the memory bank and/or whether the memory block received the request in step **2602**. Alternatively, or additionally, memory controller **2210** may directly query the internal register.

(326) If the internal register indicates that the memory bank is not loaded in an opened line (step **2606**: no), process **2600** may continue to step **2616** and a line may be loaded to a memory bank associated with the received address. In addition, memory controller **2210** or a processing unit, such as configuration manager **2212**, may signal a delay to the element requesting information

from the memory address in step **2616**. For example, if accelerator **2216** is requesting the memory information that is located an already occupied memory block, memory controller **2210** may send a delay signal to the accelerator in step **2618**. In step **2620**, configuration manager **2212** or memory controller **2210** may update the internal register to indicate a line has opened to a new memory bank or a new memory block.

(327) If the internal register indicates that the memory bank is loaded in an opened line (step **2606**: yes), process **2600** may continue to step **2608**. In step **2608**, it may be determined whether the line loaded the memory bank is being used for a different address. If the line is being used for a different address (step **2608**: yes), it would indicate that there are two instances in a single block and, therefore, they cannot be accessed simultaneously. Thus, an error or exemption signal may be sent to the element requesting information from the memory address in step **2616**. But, if the line is not being used for a different address (step **2608**: no), a line may be opened for the address and retrieve data from the target memory bank and continue to step **2614** to transmit data to the to the element requesting information from the memory address.

(328) With process **2600**, processing device **2200** has the ability to establish direct connections between processing units and the memory blocks or memory instances that contain the required information to perform a task. This organization of data would enable reading information from organized vectors in different memory instances, as well as allow the retrieval of information simultaneously from different memory blocks when a device requests a plurality of these addresses.

(329) FIG. **27** is an exemplary flowchart illustrating an execution process **2700**, consistent with disclosed embodiments. Where convenient in describing execution process **2700**, reference may be made to the identifiers of elements depicted in FIG. **22** and described above.

(330) In step **2702**, a compiler or a local unit, such as configuration manager **2212**, may receive an indication of a task that needs to be performed. The task may include a single operation (e.g., multiplication) or a more complex operation (e.g., convolution between matrixes). The task may also indicate a required computation.

(331) In step **2704**, the compiler or configuration manager **2212** may determine a number of words that is required simultaneously to perform the task. For example, configuration a compiler may determine two words are required simultaneously to perform a multiplication between vectors. In another example, a 2D convolution task, configuration manager **2212** may determine that “n” times “m” words are required for a convolution between matrices, where “n” and “m” are the matrices dimensions. Moreover, in step **2704**, configuration manager **2212** may also determine a number of cycles necessary to perform the task.

(332) In step **2706**, depending on the determinations in step **2704**, a compiler may write words that need to be accessed simultaneously in a plurality of memory banks disposed on the substrate. For instance, when a number a number of words that can be accessed simultaneously from one of the plurality of memory banks is lower than the number of words that are required simultaneously, a compiler may organize data in multiple memory banks to facilitate access to the different required words within a clock. Moreover, when configuration manager **2212** or the compiler determine a number of cycles is necessary to perform the task, the compiler may write words that are needed in sequential cycles in a single memory bank of the plurality of memory banks to prevent switching of lines between memory banks.

(333) In step **2708**, memory controller **2210** may be configured to read or grant access to at least one first word from a first memory bank from the plurality of memory banks or blocks using a first memory line.

(334) In step **2170**, a processing unit, for example one of accelerators **2216**, may process the task using the at least one first word.

(335) In step **2712**, memory controller **2210** may be configured to open a second memory line in a second memory bank. For example, based on the tasks and using the pipelined memory access approach, memory controller **2210** may be configured to open a second memory line in a second

memory block where information required for the tasks was written in step **2706**. In some embodiments, the second memory line may be opened when the task in step **2170** is about to be completed. For example, if a task requires 100 clocks, the second memory line may be opened in the 90th clock.

(336) In some embodiments, steps **2708-2712** may be executed within one line access cycle.

(337) In step **2714**, memory controller **2210** may be configured to grant access to data from at least one second word from the second memory bank using the second memory line opened in step **2710**.

(338) In step **2176**, a processing unit, for example one of accelerators **2216**, may process the task using the at least second word.

(339) In step **2718**, memory controller **2210** may be configured to open a second memory line in the first memory bank. For example, based on the tasks and using the pipelined memory access approach, memory controller **2210** may be configured to open a second memory line to the first memory block. In some embodiments, the second memory line to the first block may be opened when the task in step **2176** is about to be completed.

(340) In some embodiments, steps **2714-2718** may be executed within one line access cycle.

(341) In step **2720**, memory controller **2210** may read or grant access to at least one third word from the first memory bank from the plurality of memory banks or blocks using a second memory line in the first bank or a first line in a third bank and continuing in different memory banks.

(342) The foregoing description has been presented for purposes of illustration. It is not exhaustive and is not limited to the precise forms or embodiments disclosed. Modifications and adaptations will be apparent to those skilled in the art from consideration of the specification and practice of the disclosed embodiments. Additionally, although aspects of the disclosed embodiments are described as being stored in memory, one skilled in the art will appreciate that these aspects can also be stored on other types of computer readable media, such as secondary storage devices, for example, hard disks or CD ROM, or other forms of RAM or ROM, USB media, DVD, Blu-ray, 4K Ultra HD Blu-ray, or other optical drive media.

(343) Computer programs based on the written description and disclosed methods are within the skill of an experienced developer. The various programs or program modules can be created using any of the techniques known to one skilled in the art or can be designed in connection with existing software. For example, program sections or program modules can be designed in or by means of .Net Framework, .Net Compact Framework (and related languages, such as Visual Basic, C, etc.), Java, C++, Objective-C, HTML, HTML/AJAX combinations, XML, or HTML with included Java applets.

(344) Moreover, while illustrative embodiments have been described herein, the scope of any and all embodiments having equivalent elements, modifications, omissions, combinations (e.g., of aspects across various embodiments), adaptations and/or alterations as would be appreciated by those skilled in the art based on the present disclosure. The limitations in the claims are to be interpreted broadly based on the language employed in the claims and not limited to examples described in the present specification or during the prosecution of the application. The examples are to be construed as non-exclusive. Furthermore, the steps of the disclosed methods may be modified in any manner, including by reordering steps and/or inserting or deleting steps. It is intended, therefore, that the specification and examples be considered as illustrative only, with a true scope and spirit being indicated by the following claims and their full scope of equivalents.

Claims

1. A method performed for operating a distributed memory device comprising: determining a number of words that are required simultaneously to perform a task, the task requiring at least one computation, and providing instructions for writing words that need to be accessed simultaneously

in a plurality of memory banks when a number of words that can be accessed simultaneously from one of the plurality of memory banks is lower than the number of words that are required simultaneously; receiving, by a configuration manager, an indication to perform the task; and in response to receiving the indication, configuring a memory controller to: within a first line access cycle: access at least one first word from a first memory bank from the plurality of memory banks using a first memory line, send the at least one first word to at least one processing unit connected to the memory controller, and open a first memory line in a second memory bank to access a second address from the second memory bank from the plurality of memory banks, and within a second line access cycle: access at least one second word from the second memory bank using the first memory line, send the at least one second word to at least one processing unit connected to the memory controller, and access a third address from the first memory bank using a second memory line in the first memory bank.

2. The method of claim 1, further comprising: determining a number of cycles necessary to perform the task; and writing words that are needed in sequential cycles in a single memory bank of the plurality of memory banks.

3. The method of claim 1, further configuring a selected processing unit to transfer data to the second memory bank during the first line access cycle.

4. The method of claim 1, wherein the memory controller comprises at least two data inputs from the plurality of memory banks and at least two data outputs connected to each one of the at least one processing unit; and further comprising configuring the memory controller to: simultaneously receive data from two memory banks via the two data inputs; and simultaneously transmit data received via the two data inputs to at least one selected processing unit via the two data outputs.

5. The method of claim 1, further comprising: when a number a number of words that can be accessed simultaneously from one of the plurality of memory banks is lower than the number of words that are required simultaneously, divide the number of words required simultaneously between multiple memory banks.

6. The method of claim 1, wherein the configuration manager comprises a local memory that stores a command to be transmitted to at least one of a plurality of processing units.

7. The method of claim 1, further comprising configuring the memory controller to interrupt the task in response to receiving a request from an external interface.

8. The method of claim 1, further comprising configuring the configuration manger and the at least one processing unit to hand over access to the memory controller between each other after finalizing a task.

9. A non-transitory computer-readable medium that stores instructions that, when executed by at least one processor, cause the at least one processor to: determine a number of words that are required simultaneously to perform a task, the task requiring at least one computation; write words that need to be accessed simultaneously in a plurality of memory banks when a number of words that can be accessed simultaneously from one of the plurality of memory banks is lower than the number of words that are required simultaneously; transmit an indication to perform the task to a configuration manager; and transmit instructions to configure a memory controller to: within a first line access cycle: access at least one first word from a first memory bank from the plurality of memory banks using a first memory line, send the at least one first word to at least one processing unit connected to the memory controller, and open a first memory line in a second memory bank to access a second address from the second memory bank from the plurality of memory banks, and within a second line access cycle: access at least one second word from the second memory bank using the first memory line, send the at least one second word to at least one processing unit connected to the memory controller, and access a third address from the first memory bank using a second memory line in the first memory bank.

10. The non-transitory computer-readable medium of claim 9, wherein a selected processing unit is configured to transfer data to the second memory bank during the first line access cycle.

11. The non-transitory computer-readable medium of claim 9, wherein the memory controller comprises at least two data inputs from the plurality of memory banks and at least two data outputs connected to each one of the at least one processing unit; the memory controller is configured to simultaneously receive data from two memory banks via the two data inputs; and the memory controller is configured to simultaneously transmit data received via the two data inputs to at least one selected processing unit via the two data outputs.
 12. The non-transitory computer-readable medium of claim 9, wherein the at least one processing unit comprises a plurality of accelerators configured for pre-defined tasks.
 13. The non-transitory computer-readable medium of claim 12, wherein the plurality of accelerators comprise at least one of a vector multiply accumulate unit or a direct memory access.
 14. The non-transitory computer-readable medium of claim 9, wherein the configuration manager comprises at least one of a RISC processor or a micro-controller.
 15. The non-transitory computer-readable medium of claim 9, further comprising an external interface connected to the plurality of memory banks.
 16. The non-transitory computer-readable medium of claim 9, when a number a number of words that can be accessed simultaneously from one of the plurality of memory banks is lower than the number of words that are required simultaneously, divide the number of words required simultaneously between multiple memory banks.
 17. The non-transitory computer-readable medium of claim 9, wherein the words comprise machine instructions.
 18. The non-transitory computer-readable medium of claim 9, wherein the configuration manager comprises a local memory that stores a command to be transmitted to at least one of the at least one processing unit.
 19. The non-transitory computer-readable medium of claim 9, wherein the plurality of memory banks includes at least one of DRAM mats, DRAM, banks, flash mats, or SRAM mats.
 20. The non-transitory computer-readable medium of claim 9, wherein the at least one processing unit comprises at least one arithmetic logic unit, at least one vector handling logic unit, at least one register, and at least one direct memory access.
-