

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12405888
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Brennan; Christopher J. et al.

Dynamic performance adjustment

Abstract

A technique for operating a device is disclosed. The technique includes recording log data for the device; analyzing the log data to determine one or more performance settings adjustments to apply to the device; and applying the one or more performance settings adjustments to the device.

Inventors: Brennan; Christopher J. (Boxborough, MA), Lahiry; Akshay (Boxborough, MA)

Applicant: Advanced Micro Devices, Inc. (Santa Clara, CA)

Family ID: 89664269

Assignee: Advanced Micro Devices, Inc. (Santa Clara, CA)

Appl. No.: 18/477351

Filed: September 28, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240037031 A1	Feb. 01, 2024

Related U.S. Application Data

continuation-in-part parent-doc US 17850905 20220627 US 11860784 child-doc US 18477351

Publication Classification

Int. Cl.: G06F12/00 (20060101); G06F12/06 (20060101); G06F12/121 (20160101)

U.S. Cl.:

CPC G06F12/0646 (20130101); G06F12/121 (20130101);

Field of Classification Search

CPC: G06F (12/0802); G06F (2212/60)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5452440	12/1994	Salsburg	714/E11.197	G06F 11/3447
5579473	12/1995	Schlapp	711/119	G09G 5/39
6477314	12/2001	Tai	348/E5.053	H04N 5/9205
10049427	12/2017	Lachowsky	N/A	H04N 19/91
10339051	12/2018	Stultz	N/A	G06F 12/0646
10642764	12/2019	Gerhart	N/A	G06F 12/10
2003/0112366	12/2002	Baylon	375/E7.184	H04N 19/503
2004/0139356	12/2003	Ma	713/300	G06F 1/3215
2008/0229076	12/2007	Gonion	712/E9.034	G06F 8/443
2009/0094431	12/2008	Sarkar	711/163	G06F 12/1458
2009/0138670	12/2008	Mutlu	711/E12.001	G06F 13/1663
2013/0332681	12/2012	Miller et al.	N/A	N/A
2014/0149653	12/2013	Udipi	711/202	G06F 12/08
2015/0095472	12/2014	Cardona	709/221	G06F 9/505
2015/0139500	12/2014	Gerster	382/104	H04N 19/46
2015/0146648	12/2014	Viger et al.	N/A	N/A
2016/0283393	12/2015	Kawaba	N/A	G06F 12/0893
2017/0024568	12/2016	Pappachan	N/A	G06F 21/64
2017/0271031	12/2016	Sharon	N/A	G11C 29/52
2018/0359410	12/2017	Ain-Kedem	N/A	G06T 1/20
2019/0050160	12/2018	Dardis	N/A	G06F 3/0673
2019/0146716	12/2018	Yang et al.	N/A	N/A
2019/0172496	12/2018	Lodato	N/A	G11B 27/036
2019/0206023	12/2018	Dimitrov et al.	N/A	N/A
2019/0361808	12/2018	Subramanian	N/A	G06F 12/0802
2021/0073126	12/2020	Moyer	N/A	G06F 12/0802
2021/0240631	12/2020	Joo	N/A	G06F 12/0811
2022/0075657	12/2021	Shimada	N/A	G06F 11/3409
2022/0164299	12/2021	Moshe	N/A	G06F 21/44
2022/0269603	12/2021	Ravimohan	N/A	G06F 12/0261
2022/0374154	12/2021	Campbell	N/A	G06F 3/0622
2022/0374367	12/2021	Fang	N/A	G06F 9/30043
2023/0094834	12/2022	Leng	348/180	H04N 21/2404
2023/0096814	12/2022	Moyer	711/133	G06F 12/123
2023/0102891	12/2022	Moyer	711/118	G06F 12/127

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
1031919	12/1999	EP	N/A

Primary Examiner: Vo; Tim T

Assistant Examiner: Hasan; Mohammad S

Attorney, Agent or Firm: Volpe Koenig

Background/Summary

CROSS REFERENCE TO RELATED APPLICATIONS (1) This application is a continuation-in-part of U.S. patent application Ser. No. 17/850,905, filed on Jun. 27, 2022, which is incorporated by reference as if fully set forth herein.

BACKGROUND

(1) Varying operational settings of a computer device can have great impact on performance. Techniques for automatically varying such settings are constantly being developed.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) A more detailed understanding can be had from the following description, given by way of example in conjunction with the accompanying drawings wherein:
- (2) FIG. 1 is a block diagram of an example computing device in which one or more features of the disclosure can be implemented;
- (3) FIG. 2 is a block diagram of the device, illustrating additional details related to execution of processing tasks on the APD, according to an example;
- (4) FIG. 3 illustrates example operations performed by a logger;
- (5) FIG. 4 illustrates a dynamic runtime policy system, according to an example;
- (6) FIG. 5 illustrates operations for recording access data patterns, according to an example;
- (7) FIG. 6 illustrates example usage of the access data items, according to an example; and
- (8) FIG. 7 is a flow diagram of a method for managing memory accesses, according to an example.

DETAILED DESCRIPTION

- (9) A technique for operating a device is disclosed. The technique includes recording log data for the device; analyzing the log data to determine one or more performance settings adjustments to apply to the device; and applying the one or more performance settings adjustments to the device.
- (10) FIG. 1 is a block diagram of an example computing device **100** in which one or more features of the disclosure can be implemented. In various examples, the computing device **100** is one of, but is not limited to, for example, a computer, a gaming device, a handheld device, a set-top box, a television, a mobile phone, a tablet computer, or other computing device. The device **100** includes, without limitation, one or more processors **102**, a memory **104**, one or more auxiliary devices **106**, a storage **108**, and a last level cache (“LLC”) **110**. In some examples, the last level cache **110** is not included. An interconnect **112**, which can be a bus, a combination of buses, and/or any other communication component, communicatively links the one or more processors **102**, the memory **104**, the one or more auxiliary devices **106**, the storage **108**, and the last level cache **110**.
- (11) In various alternatives, the one or more processors **102** include a central processing unit (CPU), a graphics processing unit (GPU), a CPU and GPU located on the same die, or one or more processor cores, wherein each processor core can be a CPU, a GPU, or a neural processor. In various alternatives, at least part of the memory **104** is located on the same die as one or more of the one or more processors **102**, such as on the same chip or in an interposer arrangement, and/or at least part of the memory **104** is located separately from the one or more processors **102**. The

memory **104** includes a volatile or non-volatile memory, for example, random access memory (RAM), dynamic RAM, or a cache.

(12) The storage **108** includes a fixed or removable storage, for example, without limitation, a hard disk drive, a solid state drive, an optical disk, or a flash drive. The one or more auxiliary devices **106** include, without limitation, one or more auxiliary processors **114**, and/or one or more input/output (“IO”) devices. The auxiliary processors **114** include, without limitation, a processing unit capable of executing instructions, such as a central processing unit, graphics processing unit, parallel processing unit capable of performing compute shader operations in a single-instruction-multiple-data form, multimedia accelerators such as video encoding or decoding accelerators, or any other processor. Any auxiliary processor **114** is implementable as a programmable processor that executes instructions, a fixed function processor that processes data according to fixed hardware circuitry, a combination thereof, or any other type of processor. In some examples, the auxiliary processor(s) **114** include an accelerated processing device (“APD”) **116**.

(13) The one or more IO devices **118** include one or more input devices, such as a keyboard, a keypad, a touch screen, a touch pad, a detector, a microphone, an accelerometer, a gyroscope, a biometric scanner, or a network connection (e.g., a wireless local area network card for transmission and/or reception of wireless IEEE 802 signals), and/or one or more output devices such as a display, a speaker, a printer, a haptic feedback device, one or more lights, an antenna, or a network connection (e.g., a wireless local area network card for transmission and/or reception of wireless IEEE 802 signals).

(14) The last level cache **110** acts as a shared cache for the various components of the device **100**, such as the processor **102**, the APD **116**, and the various auxiliary devices **106**. In some implementations, there are other caches within the device **100**. For example, in some examples, the processor **102** includes a cache hierarchy including different levels such as levels 1 and 2. In some examples, each such cache level is specific to a particular logical division of the processor **102**, such as a processor core, or a processor chip, die, or package. In some examples, the hierarchy includes other types of caches as well. In various examples, one or more of the auxiliary devices **106** includes one or more caches.

(15) In some examples, the last level cache **110** is “last level” in the sense that such a cache is the last cache that the device **100** attempts to service a memory access request from before servicing that request from the memory **104** itself. For example, if a processor **102** accesses data that is not stored in any of the cache levels of the processor **102**, then the processor exports the memory access request to be satisfied by the last level cache **110**. The last level cache **110** determines whether the requested data is stored in the last level cache **110**. If the data is within the last level cache **110**, the last level cache **110** services the request by providing the requested data from the last level cache **110**. If the data is not within the last level cache **110**, the device **100** services the request from the memory **104**. As can be seen, in some implementations, the last level cache **110** acts as a final cache level before the memory **104**, which helps to reduce the overall amount of memory access latency for accesses to the memory **104**. Although techniques are described herein for operations involving the last level cache **110**, it should be understood that the techniques can alternatively be used in other types of caches or memories.

(16) FIG. 2 is a block diagram of the device **100**, illustrating additional details related to execution of processing tasks on the APD **116**, according to an example. The processor **102** maintains, in system memory **104**, one or more control logic modules for execution by the processor **102**. The control logic modules include an operating system **120**, a driver **122**, and applications **126**, and may optionally include other modules not shown. These control logic modules control various aspects of the operation of the processor **102** and the APD **116**. For example, the operating system **120** directly communicates with hardware and provides an interface to the hardware for other software executing on the processor **102**. The driver **122** controls operation of the APD **116** by, for example, providing an application programming interface (“API”) to software (e.g., applications

126) executing on the processor **102** to access various functionality of the APD **116**. The driver **122** also includes a just-in-time compiler that compiles shader code into shader programs for execution by processing components (such as the SIMD units **138** discussed in further detail below) of the APD **116**.

(17) The APD **116** executes commands and programs for selected functions, such as graphics operations and non-graphics operations, which may be suited for parallel processing. The APD **116** can be used for executing graphics pipeline operations such as pixel operations, geometric computations, and rendering an image to display device **118** based on commands received from the processor **102**. The APD **116** also executes compute processing operations that are not directly related to graphics operations, such as operations related to video, physics simulations, computational fluid dynamics, or other tasks, based on commands received from the processor **102** or that are not part of the “normal” information flow of a graphics processing pipeline, or that are completely unrelated to graphics operations (sometimes referred to as “GPGPU” or “general purpose graphics processing unit”).

(18) The APD **116** includes compute units **132** (which may collectively be referred to herein as “programmable processing units”) that include one or more SIMD units **138** that are configured to perform operations in a parallel manner according to a SIMD paradigm. The SIMD paradigm is one in which multiple processing elements share a single program control flow unit and program counter and thus execute the same program but are able to execute that program with different data. In one example, each SIMD unit **138** includes sixteen lanes, where each lane executes the same instruction at the same time as the other lanes in the SIMD unit **138** but can execute that instruction with different data. Lanes can be switched off with predication if not all lanes need to execute a given instruction. Predication can also be used to execute programs with divergent control flow. More specifically, for programs with conditional branches or other instructions where control flow is based on calculations performed by individual lanes, predication of lanes corresponding to control flow paths not currently being executed, and serial execution of different control flow paths, allows for arbitrary control flow to be followed.

(19) The basic unit of execution in compute units **132** is a work-item. Each work-item represents a single instantiation of a shader program that is to be executed in parallel in a particular lane of a wavefront. Work-items can be executed simultaneously as a “wavefront” on a single SIMD unit **138**. Multiple wavefronts may be included in a “work group,” which includes a collection of work-items designated to execute the same program. A work group can be executed by executing each of the wavefronts that make up the work group. The wavefronts may be executed sequentially on a single SIMD unit **138** or partially or fully in parallel on different SIMD units **138**. Wavefronts can be thought of as instances of parallel execution of a shader program, where each wavefront includes multiple work-items that execute simultaneously on a single SIMD unit **138** in line with the SIMD paradigm (e.g., one instruction control unit executing the same stream of instructions with multiple data). A command processor **137** is present in the compute units **132** and launches wavefronts based on work (e.g., execution tasks) that is waiting to be completed. A command processor **136** is configured to perform operations related to scheduling various wavefronts on different compute units **132** and SIMD units **138**.

(20) The parallelism afforded by the compute units **132** is suitable for graphics related operations such as pixel value calculations, vertex transformations, tessellation, geometry shading operations, and other graphics operations. A graphics processing pipeline **134** which accepts graphics processing commands from the processor **102** thus provides computation tasks to the compute units **132** for execution in parallel.

(21) The compute units **132** are also used to perform computation tasks not related to graphics or not performed as part of the “normal” operation of a graphics processing pipeline **134** (e.g., custom operations performed to supplement processing performed for operation of the graphics processing pipeline **134**). An application **126** or other software executing on the processor **102** transmits

programs (often referred to as “compute shader programs,” which may be compiled by the driver **122**) that define such computation tasks to the APD **116** for execution. Although the APD **116** is illustrated with a graphics processing pipeline **134**, the teachings of the present disclosure are also applicable for an APD **116** without a graphics processing pipeline **134**. Various entities of the APD **116**, such as the compute units **132**, can make accesses to the last level cache **110**.

(22) It is desirable to understand what memory accesses are occurring at a fine-grained level and in real time. Computer systems often provide some degree of visibility into what memory accesses are occurring and when. However, such visibility is usually limited, such as by providing only an overall count of memory accesses during a long time period. Fine-grained access information, indicating which addresses memory accesses are directed to, and what times such memory accesses are made, are usually not provided by the hardware. Such information can be very useful. For example, fine-grained memory access information can allow an application developer to identify memory access bottlenecks in an application at a very fine-grained level (e.g., to understand which accesses, made at which times, are problematic), which can allow the application developer to improve application performance by adjusting aspects of operation such as ordering of memory accesses, grouping of memory accesses, or the like. In another example, fine-grained memory access information can be provided to a runtime performance analyzer (e.g., a dynamic performance controller **404** of FIG. **4**) that can adjust execution settings based on the memory accesses.

(23) A logger **202** is illustrated in FIG. **2**. Although illustrated as part of the APD **116**, it should be understood that, in alternative examples, the logger **202** is not part of the APD **116** and is communicatively coupled to the APD **116**. The logger **202** is configured to log information that is then used by a dynamic performance controller **404** (FIG. **4**) to adjust performance settings of the device **100**.

(24) The logger **202** is configured to store performance event information into a performance log and to store memory access event information into a memory access log. The performance events include events that may or may not be directly related to memory accesses. Performance events may indicate a variety of aspects of performance, such as processing throughput, processing latency, memory access performance (e.g., amount of data successfully accessed in a given amount of time), or other aspects. In general, the performance events indicate how well the computing device **100** is performing in a given time period, and any of a variety of measures for such performance may be used. The memory access events include events such as writes to memory and reads from memory. The memory access event information indicates, either implicitly or explicitly, memory addresses associated with the memory accesses. In other words, the memory access logs store indications of what memory accesses are performed as well as what memory addresses the memory accesses are directed to. The memory address information is provided at a certain level of granularity that is not necessarily the exact address. In some examples, the memory address specifies a large chunk of memory—a memory address range—in which the memory accesses are performed, but does not provide more specific information than that. Both the performance event information and the memory access event information include or are associated with time periods or other ordering information that indicates the “time” for the associated performance events or memory access events. In some instances herein, this “time” is referred to as an epoch. In addition, the performance event information for a certain time includes memory access event references that point to associated memory access event information for that same time. These memory access event references allow correlation of performance events with memory access events, for example, during subsequent analysis or use of the performance event information and memory access event information.

(25) As stated above, the performance event information for a certain time period includes references to memory access event information for the same time period. These references allow subsequent processing to correlate performance events with memory accesses. Performance events

may indicate a variety of aspects of performance, such as processing throughput, processing latency, memory access performance (e.g., amount of data successfully accessed in a given amount of time), or other aspects. In general, the performance events indicate how well the computing device **100** is performing in a given time period, and any of a variety of measures for such performance may be used. Associating such performance events with memory access events allows subsequent analysis to first detect a performance level of interest (e.g., a drop in performance) and then to determine which memory accesses occur at that time period, and which addresses are associated with such memory accesses. In some examples, this determination allows the analysis to determine that the way in which certain memory accesses that are performed results in a particular drop in performance, or more generally to determine that some aspect of the memory accesses themselves or of processing related to the memory accesses results in, results from, or is otherwise associated with a particular drop in performance.

(26) The logger **202** is capable of recording which memory addresses the various memory accesses are targeting. In some examples, the address resolution for which such tracking occurs is variable by the logger **202**, either automatically or at the request of a different unit. The address resolution refers to the size of the memory access range for which the logger **202** stores individual items of memory access information. For example, with a 256 byte address resolution, the logger **202** stores items of memory access information for accesses within 256 byte address ranges. In an example, for a first “time,” the logger **202** detects memory accesses to addresses within a first 256 byte range and records such memory accesses as the memory access event information for the first time and for the first 256 bytes. The logger **202** thus stores information indicating that the detected number of memory accesses has occurred during the first time. No information is stored indicating which memory addresses within that 256 byte range the memory accesses occurred to. As can be seen, the address resolution indicates the specificity with which memory access events are recorded.

(27) The logger **202** records performance events on a per time basis and records memory access events on a per time basis as well. In some examples, the logger **202** records a single performance event entry for a given time and records multiple memory access event entries for a given time. Each performance event entry stores performance event information for a particular time. Each memory access event entry stores memory access event information for a particular combination of time and memory address range. In other words, in such examples, the logger **202** stores, for each time, one item of performance event information and multiple items of memory access event information, where each item of memory access event information is for a different memory address range. Each item of performance event information includes a set of performance events for a time and for multiple memory address ranges and each item of memory access event information is associated with a time and a memory address range and includes indications of which memory accesses to that memory address range occur in the associated time. It should be understood that performance event information is not specific to any memory address range and thus covers multiple memory address ranges (or can be considered to cover the entire memory address space).

(28) In some examples, the memory addresses tracked by the logger **202** are in a physical address space, as opposed to a virtual address space. A physical address space is the address space of the physical memory devices, whereas a virtual address space is the address space that is mapped to the physical address space and is thus independent from the physical addresses of the memory devices. Typically, an operating system and/or hardware (e.g., a memory controller) maintains information mapping virtual address spaces to physical address spaces. When applications or other software or hardware requests access to memory using a virtual address, the operating system and/or hardware translates such virtual addresses to physical addresses to access the memory.

(29) The address range size for the memory access event information tracked by the logger **202** may or may not be the same as the virtual address memory page size. A virtual address memory page is a portion of contiguous addresses for which a memory address translation occurs. More

specifically, typically, a virtual address has a portion (typically the most significant bits) that is considered the virtual memory page address. It is this portion that is translated to a physical address memory page. Obtaining a finer granularity physical address occurs by translating the virtual memory page address to a physical memory page address and then adding an offset. The offset is a portion of the virtual memory address other than the virtual memory page address. The size of the address range may be different from or the same as this virtual memory page size.

(30) In some examples, it is advantageous to have the tracking address range size be smaller than or equal to the virtual memory page size. This is because if the tracking address range size were larger than the virtual memory page size, then it would be possible to track in any particular item of memory access event information, information from multiple unrelated virtual memory pages, which may not be desirable. That is, tracking information from multiple unrelated virtual memory pages could result in an inability to discriminate between virtual memory pages for the aggregate states tracked in a single item of memory address tracking information. Thus, in some examples, the logger **202** limits the memory address range size to equal to or smaller than the virtual memory address page size.

(31) Above, it is stated that performance information and memory access event information is tracked for particular “times.” The term “time” is defined broadly herein. Time does not necessarily refer to wall clock time or to chip clock time or some other similar measure, although in some examples, time does refer to one such measure. Alternatively, it is possible for “time” to be measured by relation to number of tracked events occurring. In an example, “time” advances when an event that is tracked occurs. For example, when a memory access occurs, time is incremented. In another example, time is advanced for each byte of memory that is accessed over a data fabric (connection between requestor and memory) or for each byte accessed in a memory. In yet another example, time is advanced where a reference clock is advanced. Thus, in this example, in any particular item of memory access event information, a certain number of memory access events is stored. In other words, in some examples, the logger **202** tracks a certain number of events. In another example, in any particular item of performance event information, a certain number of performance events is stored.

(32) In some examples, each item of performance event information and each item of memory access event information corresponds to an “epoch.” Each epoch ends when an epoch end event occurs. In some examples, an epoch end event occurs when an item of performance event information overflows or when an item of memory access event information overflows. In some examples, an overflow occurs when the storage available in an item of performance event information or in an item of memory access event information overflows. In some examples, an overflow occurs when sufficient data has been tracked such that the amount of storage available is zero or is insufficient to store more data, for at least one item of tracked data—in other words, where there is no space left in the storage to store additional data. In some examples, the storage that overflows can track any type of information. It should be understood that this storage is the storage that stores any of the data tracked by the logger. In some examples, storage overflows when a number of performance events or memory access events is equal to a maximum for any item in a given epoch. In other examples, storage overflows when a count for statistic information for an item of performance event information or for an item of memory access event information reaches a maximum value (e.g., for an 8-bit counter, reaches 255). In some examples, when an overflow occurs, a new epoch, with new items of performance event information and memory access information, is started. It should be understood that an item of memory access event information for a given address range in a given epoch can overflow before any other item of memory access event information for a different address range in the same epoch overflows. In that instance, a new epoch is started, both for the performance event information and for all items of memory access event information. In some examples, each item of performance event information includes one or more pointers to one or more items (or all items) of memory access event information in the same

epoch as the item of performance event information. In some examples, a new epoch is not started and a new item of performance event information is created in the same epoch.

(33) FIG. 3 illustrates example operations performed by the logger **202**. The logger **202** interfaces with one or more clients **302** and one or more memories **304**. Specifically, the logger **202** monitors performance events for the client(s) **302** and memory **304**, monitors memory accesses from the clients **302** to the memory **304**, and records such events in the performance log **306** and access log **308**. The performance log **306** and access log **308** together store logged data **301**. In some example, this logged data **301** is stored in one or more memories of the APD **116** and/or the device. Logged data **301** may also be stored at other locations. In an example, the logger **202** or other element stores at least some of the logged data **301** into a memory or storage that serves as a backing store (e.g., the “backup memory” described elsewhere herein) for the logged data **301**. It should be understood that the memory **304** is any technically feasible memory, such as a cache or a memory, and such as a memory on the same chip as or on a different chip as the logger **202**, the processor **102**, or the APD **116**. In some examples, a logger **202** exists for multiple memory elements of a memory hierarchy (i.e., for each cache, each memory, and the like). In some examples, a logger **202** tracks events for multiple memories (e.g., for one or more caches and/or for one or more memories or other entities). In some such examples, each memory access log entry stores stats for multiple memories/caches. In other such examples, each memory access log entry is devoted to a single memory or cache type and each such memory access log entry includes an indication of the associated memory type.

(34) In some examples, the clients **302** are any unit of the device **100** that are capable of making memory access requests. In some examples, such units include the processor **102** and the APD **116**, as well as any other unit capable of making memory access requests, such as an auxiliary processor **114** or IO device **17**. In various examples, the logger **202** is embodied as hardware (e.g., circuitry configured to perform operations described herein), software executing on a processor, or a combination thereof. In some examples, the logger **202** is partially or fully implemented by the command processor **136**. In some examples, the command processor **136** is a programmable processor that executes software that performs some or all of the functionality of the logger **202**.

(35) As the logger **202** monitors the performance events and memory access events, the logger **202** generates performance log entries **322** and memory access log entries **324**. In some examples, the performance log entries **322** are the items of performance event information discussed above. In some examples, the memory access log entries **324** are the items of memory access event information described above. The logger **202** observes performance events and memory access events, extracts or generates information about such events, and writes such information into the performance log entries **322** and memory access log entries **324** as described elsewhere herein.

(36) Each performance log entry **322** includes a time range **326**, statistics **328**, and one or more pointers to memory access logs **330**. In some examples, the time range includes an indication of a reference clock that measures actual time (e.g., wall clock time) or that measures system clock time (e.g., number of cycles since an initial point) or that is based on system clock time for any clock of the device **100**. Note that this time range **326** is not necessarily the same type of time as the time that defines which “epoch” an entry is within, as the epoch may be based on the ordering of performance events or memory access events. Storing the time range **326** explicitly in the performance log entry **322** provides the ability to link a particular epoch to an actual point in time.

(37) The statistics **328** include information about the performance events that are tracked. Some example statistic information is for the given epoch (e.g., occurs within that epoch) and includes: the number of data fabric bytes written or read, for one or more memory types (e.g., cache, memory, or other memory type), the number of bytes written to, read from, and/or prefetched into each such memory type, the number of read instances and number of bytes involved in a compression read/modify/write operation, a number of compression metadata bytes read from or written to a memory, the number of burst writes or reads, the amount (e.g., percentage) of

bandwidth used for one or more of memory, cache, or data fabric, power management events, and user-defined events which are definable by an entity such as software. Performance log entries **322** include one or more of these items of statistic information.

(38) The number of data fabric bytes written or read includes the number of bytes written or read using a data fabric. A data fabric is the connection between clients **302** and memory **304**. Data fabrics can have a capacity (e.g., bandwidth) independent of that of the memories or clients for which the data is written or read, and thus the ability to measure the number of bytes written or read can be useful to understand. The number of bytes read or written for one or more memory types is stored on a per memory type basis. In an example a number of bytes that are read from or written to a cache in an epoch is stored in the statistics for a particular entry **322** and the number of bytes that are read from or written to a memory in the same epoch is also stored in the statistics for that entry **322**. In some examples, memory accesses are accessed to compressed data, where compression is a hardware supported operation. In some such examples, the logger **202** separately maintains the number of bytes read, written, or modified for compressed data and for uncompressed data. Thus, in some such examples, the statistics **328** for a particular entry **322** include a number of bytes read, written, or modified for uncompressed data within an epoch, and a number of bytes read, written, or modified for compressed data within the same epoch. In addition, compressed data can include or require additional metadata that specifies information required for compressing or decompressing the data or that is otherwise useful to the compression operation. In some examples, the statistics **328** store the amount of such metadata that is stored in the epoch. In some examples, memory accesses are of a burst type and a non-burst type. Generally, burst type accesses are accesses to relatively large amounts of data that is contiguous in the physical address space, whereas non-burst type accesses are accesses to individual items of data (e.g., individual words). In some examples, the statistics **328** separately store a count of burst-type accesses, in addition to a count of non-burst type accesses. The bandwidth information includes the percent (or other measure) of bandwidth capacity actually used for a memory or data fabric within the mentioned epoch.

(39) For any given performance log entry **322** for a particular epoch, the pointers to access logs **330** include pointers to the access log entries **324** for that epoch. In some examples, a single performance log entry **322** for an epoch includes pointers to all access log entries **324** in the access log **308** for that same epoch. Each access log entry **324** includes an access log entry address **332** and access log statistics **334**. The access log entry address **332** for an access log entry **324** specifies the address range of the statistics **334** for that access log entry **324**. More specifically, the entire access log **308** has an address range size (the “address resolution” described above) which specifies the granularity with which memory accesses are tracked. This address range size also indicates the range of memory addresses after the access log entry address **332** that is tracked by an access log entry **324**. In other words, each access log entry **324** tracks addresses between the access log entry address **332** and the access log entry address **332** added to the address range size. The stats **334** include information about the memory accesses tracked within the corresponding address range specified by the access log entry **332**. It should be seen that a performance log entry **322** for an epoch includes statistics **328** about multiple different address ranges that occur within an epoch, and each such address range has a different access log entry **324**, each of which includes statistics **334** about the memory accesses made to the address range within the epoch.

(40) The statistics **334** include any combination of the following, all for the epoch and the memory address range of the access log entry **324**: the number of bytes returned over the data fabric, the number of bytes written over the data fabric, the read compression ratio, the write compression ratio, the number of bytes written to or read from a memory (which such number can be stored independently for different types of memories, such as memory, caches, or other types of memories), the number of bytes prefetched into a memory, the number of bytes rinsed (where rinsing means writing dirty data back to a backing memory so that the data is no longer dirty), the

number of reads caused by compression operations, the number of atomic memory operations performed, the cache policy (including, for example, whether allocations are allowed into the cache, where allocations occur when a miss occurs, in order to store missed data into the cache, re-reference interval prediction data, which indicates the amount of “time” between re-references of a cache line), user-defined data, or any other type of information that could be stored. The read compression ratio is the ratio of the size of compressed data to the size of uncompressed data for read operations and the write compression ratio is the ratio of the size of compressed data to the size of uncompressed data for write operations. The number of reads caused by compression operations includes how many reads to data actually occur due to a compression operation. For example, reading from compressed data or writing compressed data may require reading or writing of data other than the actual data compressed, and this other data can include compression metadata or other data of a compressed block (since data may be compressed together and thus operation on one portion of compressed data may require other operations on other data that is compressed together).

(41) Log data consumers **310** are illustrated in FIG. **3** as well. These log data consumers **310** include one or more of a log data backup agent and/or a log data analyzer, either of which is embodied as hardware (e.g., circuitry), software executing on a processor, or a combination thereof. A log data backup agent stores the information from the performance log **306** and/or the access log **308** into one or more backup memories. In some examples, this storing does not, itself, trigger any logging. The one or more backup memories include one or more of the memory **104**, memories within the APD **116**, storage **108**, or other memories. The log data backup agent transfers the data from the performance log **306** and/or access log **308** when either of those logs run out of room (e.g., when a new entry is to be created but there is no space left), or in response to any other technically feasible event. The log data analyzer analyzes the performance log **306** and/or access log **308**. In some examples, the log data analyzer analyzes these logs to determine how to adjust the operation of the device **100** for better performance, and/or analyzes these logs to generate and provide conditioned information for consumption by another system or for a human (e.g., for a human developer developing an application who will use profiling data to improve performance of the application). In some examples, the log data analyzer is embodied as multiple parallel programs (such as compute shader programs executing on the APD **116**). In some examples, the log data analyzer is or is part of the dynamic performance controller **404**. For example, in some examples, the dynamic performance controller **404** includes both hardware (e.g., circuitry, a processor) as well as software executing on one or more processors such as the compute units **132** of the APD **116**. The data in the performance log **306** and access log **308** is organized to facilitate efficient parallel processing of such data. In an example, one parallel execution item (e.g., a first compute shader work-item or wavefront) analyzes a first access log entry **324** during a time period in which a second parallel execution item analyzes a second access log entry **324**. Both such parallel execution items produce a result for such analysis in parallel, which is used by a different execution item (such as a compute shader work-item or wavefront or even a thread on the processor **102**). In some examples, a different parallel execution item analyzes the performance logs **322** in parallel with the first and second parallel execution item as well. In some examples, a parallel execution item analyzes a performance log entry **322**, fetches the pointers **330**, and spawns further parallel execution items to analyze the access log entries **324** pointed to by the pointers **330**, along with information from the analysis of the performance log entries **322**. In some examples, a single execution item processes information in or derived from multiple access log entries **324**. In some examples, a single execution item processes information from different epochs. In some such examples, the single execution item aggregates information from different epochs, for example, by combining statistics for smaller address ranges and single epochs to generate statistics for larger address ranges and longer epochs. In some examples, an entity such as the dynamic performance controller **404** utilizes the information analyzed by the log data analyzer to adjust performance

settings for the APD **116** and/or other parts of the device **100**.

(42) In some examples, the logger **202** filters information, preventing that information from being written to the performance log **306** and/or the access log **308**. In some such examples, the logger **202** references data indicating the processes, virtual machines, or other entities for which logging is allowed and ignores (does not log for) accesses for which logging is not allowed. In an example, an application that has logging disabled makes memory accesses, but the logger **202** does not include information about such memory accesses in the performance log **306** and/or access log **308**.

(43) FIG. **4** illustrates a dynamic runtime policy system **400**, according to an example. The system **400** includes the APD **116**, a memory system **402**, and a dynamic performance controller **404**. In operations described herein, the dynamic performance controller **404** refers to log data **301** and generates and refers to performance policies **420** (which include access data items **506**—see FIG. **5**). The performance policies **420** are stored in any technically feasible memory, such as memory of the dynamic performance controller **404** or any other memory of the device **100**, either described herein or not described herein.

(44) In some examples, the memory system **402** is or includes the last level cache **110** of FIG. **1**. In other examples, the memory system **402** is or includes any other cache or memory that services memory access requests by the APD **116**, such as any cache or memory within the APD **116** or external to the APD **116**, or any combination thereof. In some examples, the dynamic performance controller **404** is a hardware circuit, software executing on a processor, or a combination thereof. In some examples, the dynamic performance controller **404** is included in the memory system **402** or is separate from the memory system **402**. In some examples, the dynamic performance controller **404** is included in the APD **116** or is separate from the APD **116**.

(45) The APD **116** processes data for a series of frames **406**. Each frame **406** is a full image rendered by the APD **116**. To generate graphical data over time, the APD **116** renders a series of frame, each of which includes graphical objects for a particular instant in time.

(46) To generate each frame **406**, various components of the APD **116** access data stored in the memory system **402**. Accessed data includes a wide variety of data types for rendering graphics, such as geometry data (e.g., vertex coordinates and attributes), texture data, pixel data, or a wide variety of other types of data. Aspects of performance depend in part on settings for the various elements of the memory system **402** (or for other components of the APD **116** or device **100**). The performance controller **404** is capable of, and does, dynamically apply performance policies to operations of the APD **116** and/or the device **100** based on the log data **301**. Applying these performance policies includes adjusting one or more of these settings.

(47) Some example settings include cache settings. Some examples cache settings include a no allocation setting, streaming settings, and cache aging settings. A no allocation setting includes a setting that controls whether allocations into a cache are performed as the result of a miss. A cache allocation is an event that occurs in response to a memory access that results in a miss. A memory access results in a miss in the event that the data at the address requested by the access is not in a cache. In this event, if allocation is enabled, the cache fetches the line and allocates a cache line entry in the cache and places the fetched cache line into the cache line entry. If there is a valid cache line at that entry, then that cache line must be evicted, which results in additional work. The no allocation setting prevents allocation and thus eviction if set to “yes,” and acts normally if set to “no.”

(48) Streaming settings include settings to control streaming cache fills into a cache. A streaming cache fill is a sequence of cache fills of cache lines in a contiguous address range. The settings control whether and/or how streaming occurs. Cache aging settings control how cache lines are aged. In an example, a replacement policy controls replacement of cache lines when a miss occurs. A cache controller looks for a cache line to evict based on the ages of the cache lines in a set. In an example, the cache line with the lowest age is selected for eviction. In addition, when new lines are entered into the cache, or when other events occur, the recorded age for one or more lines is

adjusted. The setting for cache aging controls how to change the age for cache lines that are tracked, thus changing which cache lines are eventually selected for eviction in the event of a miss. (49) Other settings control aspects other than cache replacement policies, such as compression settings and queue priority settings. Compression settings include whether data is compressed in a cache and control how that data is compressed (e.g., compression mode). Queue priority settings control queue priority for a memory. A memory has multiple queues that queue requests from different clients. Generally, the memory fairly acts on the queues or prioritizes one or more queues to different degrees. Acting fairly on the queues means servicing the requests in those queues in a roughly equal manner, servicing a roughly equal number of requests from each queue. Prioritizing one or more queues means servicing more requests for higher priority queues than for lower priority queues, per unit of time. Controlling these aspects means adjusting the priorities, thus adjusting the priority with which the requests of different queues are serviced. Clients to a memory can be any entity of the device **100** that can access memory.

(50) Applying a performance policy based on the log data **301** means that the dynamic performance controller **404** adjusts one or more performance settings in response to a change in one or more of the log data **301**. In various examples, any of the items of log data **301** described above (e.g., any of the described statistics **328** or statistics **334**) can be the basis for such modifications to adjustment in one or more performance settings. In an example, the dynamic performance controller **404** adjusts any one or more of cache settings or memory settings as described elsewhere herein based on one or more of the number of data fabric bytes written or read, for one or more memory types (e.g., cache, memory, or other memory type), the number of bytes written to, read from, and/or prefetched into each such memory type, the number of read instances and number of bytes involved in a compression read/modify/write operation, a number of compression metadata bytes read from or written to a memory, the number of burst writes or reads, the amount (e.g., percentage) of bandwidth used for one or more of memory, cache, or data fabric, power management events, and user-defined events which are definable by an entity such as software, the number of bytes returned over the data fabric, the number of bytes written over the data fabric, the read compression ratio, the write compression ratio, the number of bytes written to or read from a memory (which such number can be stored independently for different types of memories, such as memory, caches, or other types of memories), the number of bytes prefetched into a memory, the number of bytes rinsed (where rinsing means writing dirty data back to a backing memory so that the data is no longer dirty), the number of reads caused by compression operations, the number of atomic memory operations performed, the cache policy (including, for example, whether allocations are allowed into the cache, where allocations occur when a miss occurs, in order to store missed data into the cache, re-reference interval prediction data, which indicates the amount of “time” between re-references of a cache line), user-defined data, or any other type of information that could be stored in the log data **301**. In summary, the present disclosure contemplates a dynamic performance controller **404** that modifies one or more of the performance settings described herein based on one or more of the items of information in the log data **301**. More specifically, the present disclosure contemplates a variety of implementations of the dynamic performance controller **404**, where each such implementation is defined by a different combination of the log data **301** information described herein and the performance setting adjustment described herein, where the performance setting adjustment is made as a result of the log data information.

(51) In some examples, the dynamic performance controller **404** applies performance policies for a given frame **406** based on the log data **301** derived from operations for one or more previous frames. More specifically, patterns of memory accesses are generally very similar from frame to frame. Therefore, it is possible to make observations about operations such as memory accesses in one frame and to use those observations to control the performance policies for subsequent frames. In other words, in some examples, based on observations made about particular performance events or memory accesses in a particular frame, the dynamic performance controller **404** controls the

performance policies for the same memory accesses in one or more subsequent frames. Memory accesses are considered “the same” from frame to frame if the memory accesses occur in approximately the same place in an order of memory accesses of a frame. More specifically, in general, time-adjacent frames render very similar graphical information and thus have a very similar pattern of memory accesses. The term “approximately” allows for the possibility that some memory accesses that occur in one frame may not occur in another frame, since the frames can be somewhat different, meaning that the order of memory accesses between frames will not be identical.

(52) An example of a manner in which the dynamic performance controller **404** controls performance policies in response to log data **301** is now provided. In this example, the dynamic performance controller **404** controls, in response to the log data **301**, a cache access policy that controls three aspects: the age at which a new cache line is brought into the cache (e.g., as the result of a miss), the manner in which ages of the cache lines are updated when a miss occurs (e.g., how cache lines other than the accessed cache line are aged), and the age that a cache line is updated to when a hit occurs for the cache line (e.g., what age is the cache line set to when it is accessed).

(53) It should be noted that the age referred to is an age in a least recently used cache replacement policy. More specifically, in such a policy, a cache is a set-associative cache that includes a number of sets each having a plurality of ways. Any given cache line is mapped to at least one set and can be placed into any way in that set, but cannot be placed in any set that the cache line is not mapped to. A cache line is defined as a portion of the memory address space. In an example, each cache line has a cache line address which includes a certain number of the most significant bits of the address, such that the remaining least significant bits specify an offset in that cache line. Typically, though not necessarily, a cache line is the minimum unit of data read into or written back from the cache. When a cache line is to be brought into a set of the cache and there are no free ways for the cache line, the cache evicts one of the cache lines in the set. The replacement algorithm selects the oldest cache line as the cache line to evict based on the ages of the cache lines in that set. The cache also places the new cache line into the cache with a “starting age.” In the event of a hit, the cache updates the age of the cache line for which the hit occurred based on the memory access policy. Additional details are provided elsewhere herein.

(54) As described above, the log data **301** includes information about memory accesses made to particular memory address ranges and at particular “times.”

(55) In one example, the dynamic performance controller **404** determines the re-reference interval and the re-reference intensity for the accesses corresponding to the data in the log data **301**. Then, the dynamic performance controller **404** determines a policy according to the re-reference interval and re-reference intensity for the accesses corresponding to the information in the log data **301** that is analyzed. In some examples, the dynamic performance controller **404** determines the policy based on the re-reference interval and re-reference intensity from multiple frames. In some examples, the dynamic performance controller **404** remembers this policy for use in the subsequent frame.

(56) FIG. 5 illustrates operations for recording access data patterns **506**, according to an example. FIG. 5 illustrates a single frame—frame 1 **508(1)** in which a number of accesses **512** are occurring. In this example, a first set of accesses—accesses **512(1)** to **512(4)** are to a first address range AR1 and occur within a certain time period. The logger **202** identifies the three sets of accesses **512** and generates information for the log data **301**. The dynamic performance controller **404** generates access data items **506** based on the log data **301**, where the access data items **506** correspond to each of the sets of accesses **512**. As described above, the log data **301** includes one or both of the performance log entries **322** and/or the memory access log entries **324** described with respect to FIG. 3. The access data items **506** include information regarding how to “condition” subsequent memory accesses. In other words, the dynamic performance controller **404** makes a determination based on the log data **301** regarding how to “condition” subsequent data accesses in subsequent

frames. The dynamic performance controller **404** records such determinations as the access data items **506**.

(57) FIG. **6** illustrates example usage of the access data items **506**, according to an example. The dynamic performance controller **404** examines the access data items **506** and “conditions” accesses made in frames subsequent to the frame(s) in which the access data items **506** are recorded.

“Conditioning” memory accesses means setting parameters for the memory accesses, where the parameters correspond to the determined re-reference interval values and the determined re-reference intensity values. For example, in a first frame, the dynamic performance controller **404** records a first access data item **506** for a set of memory accesses that access the same page. The first access data item **506** indicates a certain policy for that set of memory accesses. In a subsequent frame, the dynamic performance controller **404** identifies the same accesses (that is, the accesses of the first access data item **506**) and “conditions” those accesses according to the policy stored in the first access data item **506**. More specifically, the dynamic performance controller **404** configures those accesses to use the policy that corresponds to the first access data item **506**. Such a policy indicates that these accesses in the subsequent frame should be made according to a certain set of settings. In some examples, the policy indicates the following: how to age cache lines other than the accessed cache line and in the same set as the cache line, in the event of a miss; what age to insert a new cache line into the cache line; and what age to set a cache line to when a hit occurs.

(58) Aging cache lines in the event of a miss is now described. When a miss occurs in a cache, the cache identifies the cache line to bring into the cache. The cache also determines which set the identified cache line is to be brought in to. If any cache line in that set has an age that is equal to or greater than a threshold (e.g., 3 if the age counter is two bits), then the cache selects that cache line for eviction. If no cache lines have an age that is higher than a threshold, then the cache ages the cache lines of the set. The setting of how cache lines are aged in the event of a miss indicates which cache lines of a set are aged in the event that a miss occurs in that set and no cache line has an age that is equal to or greater than a threshold. In some examples, this setting indicates, by age, which cache lines will be aged in the event of a miss. In other words, the setting indicates the ages of cache lines that will be aged in the event of a miss. In an example of the setting, the setting indicates that cache lines of all ages lower than the above threshold (the “eviction threshold”) are aged. In another example, the setting indicates that cache lines having ages above an age trigger threshold, which can be lower than the eviction threshold, are aged. In this situation, cache lines with ages lower than or equal to the low threshold are not aged in the event of a miss. In sum, in some examples, the re-reference interval and re-reference intensity of a set of memory accesses for one frame indicates how to age cache lines (specifically, the ages of cache lines that will be aged) in the event of a miss for the set of memory accesses in a subsequent frame.

(59) The setting of what age to insert a new cache line into the cache is now described. When a cache line is brought into the cache as the result of a miss, the cache line is initially given a particular age. In some examples, this “starting” age is the lowest possible age, some intermediate age that is above the lowest possible age, or the maximum age. Again, this setting is dependent on the memory accesses reflected in an access data item **506**. Thus, the access data item **506**, corresponding to a memory page, indicates the starting age for a cache line in the event that a miss occurs and the cache line is copied into the cache.

(60) What age to set a new cache line to when a hit occurs is now described. When a hit occurs to a cache line, the cache updates the age of that cache line (e.g., to indicate that the cache line is “newer”). In some examples, this setting indicates that the cache line is to have a particular age (such as 0) when a hit occurs for that cache line. In other examples, the setting indicates that the cache line is to have a different age such as 1, 2, or 3 (for a two bit age counter) in the event of a hit. In other examples, this setting indicates that the age of a cache line is to be modified in a particular manner, such as by decrementing the age by a number like 1. In sum, the access data item **506**, corresponding to a memory page, indicates how the age of a cache line is modified when

a hit occurs for that cache line.

(61) It should be understood that conditioning a particular cache access according to a policy means causing the cache access to occur with that policy. For the insertion age policy, this policy is applied for an access conditioned according to that policy if that access results in a miss. The cache line that is to be brought in is brought in according to the age specified by the policy. For the aging policy, this occurs for the access conditioned according to the policy in the event that that access results in a miss. In this situation, the aging policy causes the other cache lines for the same set to be aged as specified by the policy. For the policy that defines what age a cache line will be set to in the event of a hit, when an access occurs that is conditioned according to a policy, and that access results in a hit, the policy causes the age of the cache line targeted by the access to be set according to the policy.

(62) In FIG. 6, three sets **602** of data accesses **612** are illustrated. The dynamic performance controller **404** conditions these accesses according to the access data items **506** created in frame 1 **508(1)** of FIG. 5. For the data accesses **612** of the first set **602**, the dynamic cache policy controller **304** identifies the first access data item **506(1)**. This access data item **506** indicates a certain manner in which to condition the accesses **612** of the set **602(1)**. As described elsewhere herein, the access data item **506** indicates a policy, which, in an example, indicates one or more of how to age cache lines when a miss occurs, what age to set new cache lines to, and what age to set cache lines to in the event of a hit. Additionally, this policy is dependent on the data recorded in frame 1 **508(1)** about the same accesses from that frame. The dynamic cache policy controller **304** causes the policy applied to the accesses of a given set **602** to be applied based on the recorded access data item **506** for the same accesses in a previous frame. Similar activity is applied for set **602(2)** and set **602(3)**.

(63) In some examples, the dynamic performance controller **404** determines which accesses in a particular frame are “the same as” accesses in a previous frame in the following manner. In each frame, memory accesses occur in a particular sequence. This sequence is mostly repeated from frame to frame. Thus, the dynamic cache policy controller **304** tracks the sequence of memory accesses to identify which accesses are “the same” from frame to frame. It is true that some accesses may differ from frame to frame, so the dynamic cache policy controller **304** accounts for such differences. For example, the dynamic cache policy controller **304** may notice omitted accesses, added accesses, or other modifications, and account for those. In some examples, memory accesses that occur to the same page and at the same “time” in different frames are considered to be “the same” memory accesses. In examples, the “time” is defined based on access order. For example, the first through one hundredth accesses are in a first “time,” the one hundred first through two hundredth are in a second “time,” and so on. In some examples, the time represents the time range **326** information stored in the log data **301**. In some examples, the dynamic performance controller **404** determines that accesses occur at the same time in two different frames by applying a start-of-frame adjustment to the time range **326** of the memory accesses. For example, by subtracting the time of the start of the frame from the time for a given access from that frame, the dynamic performance controller **404** obtains the time within that frame associated with the accesses. Thus, the dynamic performance controller **404** is able to equate accesses from different frames to each other. In an additional example, the dynamic performance controller **404** considers the order of accesses and also applies a fudge factor sliding window that allows for memory accesses that occur at somewhat different times in different frames to be considered the same, accounting for the fact that operations in different frames may not take the exact same amount of time. The dynamic performance controller **404** also considers the addresses accessed. If the addresses are the same and are the same and the times are also considered the same, then the dynamic performance controller **404** considers the accesses to be the same accesses in different frames. If either the addresses are not the same or the times are not considered the same, then the dynamic performance controller **404** considers the accesses to not be the same accesses in different

frames. Any other technically feasible means for determining the time is possible.

(64) It should be understood that the manner in which accesses are “conditioned” is based on the access data items **506** recorded in a previous frame. Thus, for one particular access data item **506**, which indicates a particular combination of re-reference interval and re-reference intensity, corresponding accesses are made with a first set of parameters including manner of aging cache lines, age to insert new cache lines, and age to update hit cache lines. For another particular access data item **506**, which indicates a different combination of re-reference interval and re-reference intensity, corresponding accesses are made with a second set of parameters including manner of aging cache lines, age to insert new cache lines, and age to update hit cache lines. At least one of the parameters of the second set is different from at least one of the parameters of the first set.

(65) FIG. 7 is a flow diagram of a method **700** for managing performance settings in accordance with logged data, according to an example. Although described with respect to the system of FIGS. 1-6, those of skill in the art will understand that any system, configured to perform the steps of the method **700** in any technically feasible order, falls within the scope of the present disclosure.

(66) At step **702**, the logger **202** records log data **301** for the device **100**. This recording occurs as described elsewhere herein. For example, the logger **202** records any of the information described above, including the statistics **328** for the performance log **306** and the statistics **334** for the access log **308**. The logger **202** also records times for this information, again as described elsewhere herein, such as in the time range **326**.

(67) At step **704**, a dynamic performance controller **404** analyzes the log data **301** to determine performance setting adjustments. As described elsewhere herein, the present disclosure contemplates a variety of implementations of the dynamic performance controller **404** that considers any combination of the information stored in the log data **301** to make adjustments to a variety of settings as described elsewhere herein.

(68) At step **706**, the dynamic performance controller **404** applies the performance setting adjustments to operation of the device **100**. In an example, the dynamic performance controller **404** conditions memory accesses according to the performance setting adjustments, which causes the memory accesses to be performed according to those performance settings adjustments.

(69) In some examples, the performance settings applied in step **706** are applied to a frame that is subsequent to the frame in which the log data **301** is recorded at step **702**. In other words, in some examples, the dynamic performance controller **404** analyzes log data **301** obtained at one frame to generate performance settings adjustments to apply to a subsequent frame. In some examples, the dynamic performance controller **404** applies settings adjustments determined for memory accesses at a particular time in one frame to the “same” memory accesses applied at approximately the same time in a subsequent frame. In some examples, memory accesses are the same in different frames if the memory accesses are to the same addresses and occur at approximately the same time from the beginning of the frame. As described elsewhere herein, memory accesses occur at approximately the same time if the recorded time for the memory accesses minus the frame start time for the frame in which those accesses were made is approximately the same. In some examples, being approximately the same means being the same to within a threshold (e.g., an absolute threshold or a percentage threshold).

(70) The elements in the figures are embodied as, where appropriate, software executing on a processor, a fixed-function processor, a programmable processor, or a combination thereof. The processor **102**, last level cache **110**, interconnect **112**, memory **104**, storage **108**, various auxiliary devices **106**, APD **116** and elements thereof, logger **202**, clients **302**, and the dynamic performance controller **404**, include at least some hardware circuitry and, in some implementations, include software executing on a processor within that component or within another component.

(71) It should be understood that many variations are possible based on the disclosure herein. Although features and elements are described above in particular combinations, each feature or element can be used alone without the other features and elements or in various combinations with

or without other features and elements.

(72) The methods provided can be implemented in a general purpose computer, a processor, or a processor core. Suitable processors include, by way of example, a general purpose processor, a special purpose processor, a conventional processor, a digital signal processor (DSP), a plurality of microprocessors, one or more microprocessors in association with a DSP core, a controller, a microcontroller, Application Specific Integrated Circuits (ASICs), Field Programmable Gate Arrays (FPGAs) circuits, any other type of integrated circuit (IC), and/or a state machine. Such processors can be manufactured by configuring a manufacturing process using the results of processed hardware description language (HDL) instructions and other intermediary data including netlists (such instructions capable of being stored on a computer readable media). The results of such processing can be maskworks that are then used in a semiconductor manufacturing process to manufacture a processor which implements features of the disclosure.

(73) The methods or flow charts provided herein can be implemented in a computer program, software, or firmware incorporated in a non-transitory computer-readable storage medium for execution by a general purpose computer or a processor. Examples of non-transitory computer-readable storage mediums include a read only memory (ROM), a random access memory (RAM), a register, cache memory, semiconductor memory devices, magnetic media such as internal hard disks and removable disks, magneto-optical media, and optical media such as CD-ROM disks, and digital versatile disks (DVDs).

Claims

1. A method for operating a device, the method comprising: recording log data for the device, wherein the recording occurs in a first graphics frame; analyzing the log data to determine one or more performance settings adjustments to apply to the device; and applying the one or more performance settings adjustments to the device, wherein the applying occurs for operations of a second frame, wherein the second frame is subsequent to the first frame.
2. The method of claim 1, wherein the operations of the second frame comprise memory accesses that occur within a threshold amount of time in the second frame as memory accesses for which the log data is recorded occur in the first frame.
3. The method of claim 1, wherein the log data includes one or both of time range data and statistics data.
4. The method of claim 2, wherein the log data comprise one or both of performance statistics and memory access statistics.
5. The method of claim 1, wherein the one or more performance settings comprise one or more of cache settings and memory settings.
6. The method of claim 1, wherein the one or more performance settings comprise one or more of a no allocation setting, a streaming setting, a cache aging setting, a compression setting, and a queue priority setting.
7. The method of claim 1, wherein analyzing the log data to determine one or more performance settings to apply to the device comprises determining that the one or more performance settings should change in response to a change in the log data.
8. The method of claim 1, wherein performance log entries of the log data include references to memory access log entries.
9. A system, comprising: a processor; a cache; and circuitry configured to: record log data for the system, wherein the recording occurs in a first graphics frame; analyze the log data to determine one or more performance settings adjustments to apply to the system; and apply the one or more performance settings to the system, wherein the applying occurs for operations of a second frame, wherein the second frame is subsequent to the first frame.
10. The system of claim 9, wherein the operations of the second frame comprise memory accesses

that occur within a threshold amount of time in the second frame as memory accesses for which the log data is recorded occur in the first frame.

11. The system of claim 9, wherein the log data includes one or both of time range data and statistics data.

12. The system of claim 10, wherein the log data comprise one or both of performance statistics and memory access statistics.

13. The system of claim 9, wherein the one or more performance settings comprise one or more of cache settings and memory settings.

14. The system of claim 9, wherein the one or more performance settings comprise one or more of a no allocation setting, a streaming setting, a cache aging setting, a compression setting, and a queue priority setting.

15. The system of claim 9, wherein analyzing the log data to determine one or more performance settings to apply to the system comprises determining that the one or more performance settings should change in response to a change in the log data.

16. The system of claim 9, wherein performance log entries of the log data include references to memory access log entries.

17. A non-transitory computer-readable medium storing instructions that, when executed by a processor, cause the processor to perform operations comprising: recording log data for a device, wherein the recording occurs in a first graphics frame; analyzing the log data to determine one or more performance settings adjustments to apply to the device; and applying the one or more performance settings adjustments to the device, wherein the applying occurs for operations of a second frame, wherein the second frame is subsequent to the first frame.

18. The non-transitory computer-readable medium of claim 17, wherein the operations of the second frame comprise memory accesses that occur within a threshold amount of time in the second frame as memory accesses for which the log data is recorded occur in the first frame.

19. The non-transitory computer-readable medium of claim 17, wherein the log data includes one or both of time range data and statistics data.

20. The non-transitory computer-readable medium of claim 19, wherein the log data comprise one or both of performance statistics and memory access statistics.
