

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12405934
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Schillace; Samuel Edward et al.

Storing entries in and retrieving information from an embedding object memory

Abstract

Methods, systems, and media for storing entries in and/or retrieving information from an embedding object memory are provided. In some examples, a content item is received that has content data. The content data associated with the content item may be provided to one or more semantic embedding models that generate semantic embeddings. From one or more of the semantic embedding models, one or semantic embeddings may be received. The one or more semantic embeddings may then be inserted into the embedding object memory. The semantic embeddings may be associated with respective indications corresponding to a reference to source data associated with the semantic embeddings. Further, the insertion may trigger a spatial storage operation to store a vector representation of the one or more semantic embeddings. A plurality of collections of stored embeddings may be received from the embedding object memory, based on a provided input, to determine an action.

Inventors: Schillace; Samuel Edward (Portola Valley, CA), Madan; Umesh (Bellevue, WA), Lucato; Devis (Kirkland, WA)

Applicant: Microsoft Technology Licensing, LLC (Redmond, WA)

Family ID: 91473806

Assignee: Microsoft Technology Licensing, LLC (Redmond, WA)

Appl. No.: 18/122563

Filed: March 16, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240202173 A1	Jun. 20, 2024

Related U.S. Application Data

Publication Classification**Int. Cl.:** G06F16/21 (20190101); G06F16/245 (20190101); G06F16/2457 (20190101)**U.S. Cl.:****CPC** G06F16/219 (20190101); G06F16/24573 (20190101);**Field of Classification Search****USPC:** None

References Cited**U.S. PATENT DOCUMENTS**

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
10171656	12/2018	Pullamplavil	N/A	N/A
10950231	12/2020	Kockerbeck	N/A	N/A
11017780	12/2020	Steelberg	N/A	N/A
11095468	12/2020	Pandey et al.	N/A	N/A
11106736	12/2020	Newman	N/A	G06F 16/284
11328368	12/2021	Labrie	N/A	N/A
11443164	12/2021	Dalli	N/A	N/A
11573993	12/2022	Nelson et al.	N/A	N/A
11605387	12/2022	Muralitharan	N/A	N/A
11823477	12/2022	Ramezani	N/A	G06F 16/24
2003/0115548	12/2002	Melgar	N/A	N/A
2014/0040262	12/2013	Winter	N/A	N/A
2016/0070911	12/2015	Okereke	N/A	N/A
2017/0212886	12/2016	Sarikaya et al.	N/A	N/A
2018/0060334	12/2017	Jensen	N/A	G06F 16/3344
2018/0151081	12/2017	Chen	N/A	G06F 16/3347
2019/0028520	12/2018	Nawrocki	N/A	N/A
2019/0384813	12/2018	Mohamed	N/A	N/A
2020/0133071	12/2019	Dai	N/A	N/A
2020/0311122	12/2019	Ramamurthy et al.	N/A	N/A
2020/0311162	12/2019	Xu	N/A	G06F 16/9536
2020/0403817	12/2019	Daredia et al.	N/A	N/A
2021/0109769	12/2020	Chenguang	N/A	N/A
2021/0117479	12/2020	Liu	N/A	N/A
2021/0182341	12/2020	Mullins	N/A	N/A
2021/0224336	12/2020	Bright	N/A	N/A
2021/0271707	12/2020	Lin	N/A	G06F 16/538
2021/0303638	12/2020	Zhong	N/A	G06F 16/9535
2021/0342711	12/2020	Mokeyev	N/A	G06F 16/3347
2022/0100831	12/2021	Moreno	N/A	G06F 16/21

2022/0101861	12/2021	Antos	N/A	N/A
2022/0198156	12/2021	Rao	N/A	N/A
2022/0199079	12/2021	Hanson	N/A	N/A
2022/0200934	12/2021	Dutta	N/A	N/A
2022/0342900	12/2021	Basu	N/A	N/A
2022/0358713	12/2021	Krishnamurthy	N/A	N/A
2022/0385758	12/2021	Tadesse et al.	N/A	N/A
2023/0035076	12/2022	Wang	N/A	N/A
2023/0111517	12/2022	Anderson	N/A	N/A
2023/0115420	12/2022	Dabas	N/A	N/A
2023/0135179	12/2022	Mielke	N/A	N/A
2023/0139347	12/2022	Bondarenko	N/A	N/A
2023/0142718	12/2022	Asgekar	N/A	G06F 16/90335
2023/0153641	12/2022	Manda et al.	N/A	N/A
2023/0153700	12/2022	Lindgren	N/A	G06F 16/245
2023/0155903	12/2022	Song	N/A	N/A
2023/0386469	12/2022	Horton	N/A	N/A
2024/0078376	12/2023	Li	N/A	G06F 16/583
2024/0136070	12/2023	El Saadawi	N/A	G06F 16/90335
2024/0201959	12/2023	Callegari	N/A	N/A
2024/0202215	12/2023	Shashanka	N/A	G06F 16/285
2024/0202451	12/2023	Schillace	N/A	N/A
2024/0202452	12/2023	Schillace	N/A	N/A
2024/0202460	12/2023	Schillace	N/A	N/A
2024/0202582	12/2023	Schillace	N/A	N/A
2024/0202584	12/2023	Schillace	N/A	N/A
2024/0205037	12/2023	Callegari	N/A	N/A

OTHER PUBLICATIONS

International Search Report and Written Opinion received for PCT Application No.

PCT/US2023/081290, Mar. 11, 2024, 17 pages. cited by applicant

International Search Report and Written Opinion received for PCT Application No.

PCT/US2023/082379, Apr. 22, 2024, 12 pages. cited by applicant

“Long-term Memory for AI,” Vector Database for Vector Search, Pinecone Systems, Inc., retrieved from: <https://www.pinecone.io/>, Mar. 27, 2023, 10 pages. cited by applicant

Kan, “Not All Vector Databases Are Made Equal,” Towards Data Science, Oct. 2, 2021, retrieved from: <https://towardsdatascience.com/milvus-pinecone-vespa-weaviate-vald-gsi-what-unites-these-buzz-words-and-what-makes-each-9c65a3bd0696>, Mar. 27, 2023, 15 pages. cited by applicant

Notice of Allowance mailed on Aug. 27, 2024, in U.S. Appl. No. 18/129,758, 12 pages. cited by applicant

Notice of Allowability mailed on Dec. 4, 2024, in U.S. Appl. No. 18/129,758, 5 pages. cited by applicant

Anonymous: “Is there way to submit a table as an JSON file?”, Aug. 19, 2022, Retrieved from the Internet: URL: <https://web.archive.org/web/20220819203852/https://community.openai.com/t/is-there-way-to-submit-a-table-as-an-json-file/5299>, 1 page. cited by applicant

Anonymous: “OpenAI—Spreadsheet creator”, Aug. 19, 2022, XP093140385, Retrieved from the Internet: URL: <https://platform.openai.com/examples/default-spreadsheet-gen>, 1 page. cited by applicant

Anonymous: “Skill Chaining—Wikipedia, the free encyclopedia”, Nov. 27, 2015, Retrieved from the Internet: URL:

https://web.archive.org/web/20151127070537/https://en.wikipedia.org/wiki/Skill_chaining, 1 page. cited by applicant

Beaumont, Romain., “Semantic search with embeddings: index anything”, Dec. 1, 2020, Retrieved from the Internet: URL: <https://web.archive.org/web/20221212165434/https://rom1504.medium.com/semantic-search-with-embeddings-index-anything-8fb18556443c>, retrieved on Jan. 26, 2024, 27 pages. cited by applicant

Chan Irene, “How to Create Meetings Summaries with OpenAI GPT-3”, Dec. 16, 2022, 04 pages. cited by applicant

Fukamizu, et al., “Generation High resolution 3D model from natural language by Generative Adversarial Network”, arxiv.org, Cornell University Library, 201 Olin Library Cornell University Ithaca, NY 14853, Jan. 22, 2019, 9 pages. cited by applicant

Holovaty, Adrian: “Adventures in generating music via ChatGPT text prompts”, Dec. 1, 2022, Retrieved from the Internet: URL: <https://web.archive.org/web/20221217201750/https://www.holovaty.com/writing/chatgpt-music-generation/>, 10 pages. cited by applicant

International Search Report and Written Opinion received for PCT Application No. PCT/US2023/081254, mailed on Apr. 10, 2024, 10 pages. cited by applicant

International Search Report and Written Opinion received for PCT Application No. PCT/US2023/081261, mailed on Feb. 12, 2024, 15 pages. cited by applicant

International Search Report and Written Opinion received for PCT Application No. PCT/US2023/081271, mailed on Mar. 13, 2024, 12 pages. cited by applicant

International Search Report and Written Opinion received for PCT Application No. PCT/US2023/081303, mailed on Mar. 26, 2024, 15 pages. cited by applicant

International Search Report and Written Opinion received for PCT Application No. PCT/US2023/081323, mailed on Mar. 13, 2024, 17 pages. cited by applicant

International Search Report and Written Opinion received for PCT Application No. PCT/US2023/081342, mailed on Mar. 26, 2024, 14 pages. cited by applicant

Konidaris, et al., “Skill Discovery in Continuous Reinforcement Learning Domains using Skill Chaining”, Dec. 10, 2009, Retrieved from the Internet: URL: https://proceedings.neurips.cc/paper_files/paper/2009/file/e0cf1f47118daebc5b16269099ad7347-Paper.pdf, 9 pages. cited by applicant

Magnani, et al., “Semantic Retrieval at Walmart”, Proceedings of the 59th ACM/IEEE Design Automation Conference, ACM/PUB27, New York, NY, USA, Aug. 14, 2022, pp. 3495-3503. cited by applicant

Non-Final Office Action mailed on May 8, 2024, in U.S. Appl. No. 18/129,758, 21 pages. cited by applicant

T, Vincent., “GPT-3 and Code Generation—AI-enabled Instant Software Development”, Jan. 12, 2021, retrieved from the Internet: URL: <https://becominghuman.ai/gpt-3-and-code-generation-ai-enabled-instant-software-development-270795077cbd>, 10 pages. cited by applicant

Non-Final Office Action mailed on Jan. 16, 2025, in U.S. Appl. No. 18/129,668, 22 pages. cited by applicant

Non-Final Office Action mailed on Mar. 20, 2025, in U.S. Appl. No. 18/129,697, 08 pages. cited by applicant

Notice of Allowance mailed on Feb. 12, 2025, in U.S. Appl. No. 18/129,758, 05 pages. cited by applicant

Non-Final Office Action mailed on Apr. 15, 2025, in U.S. Appl. No. 18/129,783, 40 pages. cited by applicant

Non-Final Office Action mailed on Apr. 18, 2025, in U.S. Appl. No. 18/129,772, 15 pages. cited by applicant

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION (1) This application claims priority to U.S. Provisional Application No. 63/433,619, titled “STORING ENTRIES IN AND RETRIEVING INFORMATION FROM AN EMBEDDING OBJECT MEMORY,” filed on Dec. 19, 2022, the entire disclosure of which is hereby incorporated by reference in its entirety.

BACKGROUND

(1) Computing devices may be relied on to perform any of a variety of different tasks. Further, computing devices may receive large quantities of content information, such as from video inputs, audio inputs, data transmissions, applications being executed, etc. Some systems may categorize and store the large quantities of content information that computing devices receive to compare related content objects for further processing. For example, the systems may use keyword searches to iterate through each content object from the large quantities of content information to find which content objects are related to the keyword. However, keyword searches may not consider the abstract meaning behind content objects, and may therefore be relatively inaccurate. Further, storing and retrieving objects from large quantities of content information may be computationally inefficient, such as by requiring a relatively large amount of memory to store the content information that is to be searched upon.

(2) It is with respect to these and other general considerations that embodiments have been described. Also, although relatively specific problems have been discussed, it should be understood that the embodiments should not be limited to solving the specific problems identified in the background.

SUMMARY

(3) Aspects of the present disclosure relate to methods, systems, and media for storing entries in and/or retrieving information from an embedding object memory.

(4) In some examples, a content item (e.g., a document, a skill, data object, etc.) is received that has content data. The content data associated with the content item may be provided to one or more semantic embedding models that generate semantic embeddings. From one or more of the semantic embedding models, one or semantic embeddings may be received. The one or more semantic embeddings may then be inserted into the embedding object memory. The semantic embeddings may be associated with respective indications corresponding to a reference to source data associated with the semantic embeddings. Further, the insertion may trigger a spatial storage operation to store a vector representation of the one or more semantic embeddings. After the embeddings are stored, a plurality of collections of stored embeddings may be received from the embedding object memory, such as based on a provided input. Further, an action may be determined based on a subset of the collections of stored embeddings and the provided input. The subset of the collections of stored embeddings may be retrieved based on a similarity to the input embedding.

(5) This summary is provided to introduce a selection of concepts in a simplified form that are further described below in the Detailed Description. This summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used to limit the scope of the claimed subject matter. Additional aspects, features, and/or advantages of examples will be set forth in part in the following description and, in part, will be apparent from the description, or may be learned by practice of the disclosure.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) Non-limiting and non-exhaustive examples are described with reference to the following Figures.
- (2) FIG. 1 illustrates an overview of an example system according to some aspects described herein.
- (3) FIG. 2 illustrates examples of private content and public content, according to some aspects described herein.
- (4) FIG. 3 illustrates an example flow of storing entries in an embedding object memory, according to some aspects described herein.
- (5) FIG. 4 illustrates an example method for storing entries in an embedding object memory, according to some aspects described herein.
- (6) FIG. 5 illustrates an example flow for retrieving information from an embedding object memory, according to some aspects described herein.
- (7) FIG. 6 illustrates an example vector space, according to some aspects described herein.
- (8) FIG. 7 illustrates an example method for retrieving information from an embedding object memory, according to some aspects described herein.
- (9) FIGS. 8A and 8B illustrate overviews of an example generative machine learning model that may be used according to aspects described herein.
- (10) FIG. 9 illustrates a block diagram illustrating example physical components of a computing device with which aspects of the disclosure may be practiced.
- (11) FIG. 10 illustrates a simplified block diagrams of a computing device with which aspects of the present disclosure may be practiced.
- (12) FIG. 11 is a simplified block diagram of a distributed computing system in which aspects of the present disclosure may be practiced.

DETAILED DESCRIPTION

- (13) In the following detailed description, references are made to the accompanying drawings that form a part hereof, and in which are shown by way of illustrations specific embodiments or examples. These aspects may be combined, other aspects may be utilized, and structural changes may be made without departing from the present disclosure. Embodiments may be practiced as methods, systems or devices. Accordingly, embodiments may take the form of a hardware implementation, an entirely software implementation, or an implementation combining software and hardware aspects. The following detailed description is therefore not to be taken in a limiting sense, and the scope of the present disclosure is defined by the appended claims and their equivalents.
- (14) As mentioned above, computing devices may be relied on to perform any of a variety of different tasks. Further, computing devices may receive large quantities of content information, such as from video inputs, audio inputs, data transmissions, applications being executed, etc. Some systems may categorize and store the large quantities of content information that computing devices receive to compare related content objects for further processing. For example, the systems may use keyword searches to iterate through each content object from the large quantities of content information to find which content objects are related to the keyword. However, keyword searches may not consider the abstract meaning behind content objects, and may therefore be relatively inaccurate. Further, storing and retrieving objects from large quantities of content information may be computationally inefficient, such as by requiring a relatively large amount of memory to store the content information that is to be searched upon.
- (15) The content information discussed herein may be unstructured or semi-structured content, such as raw text or images. Standard Indexing techniques may be otherwise used for structured

content, such as databases. Systems and methods provided herein may be used for text and other unstructured content, which despite integration of statistics may fundamentally remain a bag of words solutions (e.g., requiring instances of features associated with content to be individually identified, selected, and/or counted).

(16) When building an operating system and/or kernel for semantic models (e.g., semantic embedding models) to perform actions in a system, the system may have the ability to learn skills and/or commands, such as by referencing external sources. The system may further have an embedding memory outside of a semantic model, such that semantic context can be built (e.g., in real-time). Further, the system may generate, launch and execute actions based on the embedding memory.

(17) Recent advancements in artificial intelligence and machine-learning techniques allow for abstract meaning to be extracted from content objects in the form of embeddings that can be mathematically compared against the abstract meaning of other content objects in the form of embeddings, such that a quantitative similarity between the embeddings, and by consequence the content objects from which the embeddings were generated, can be determined. A similarity between content objects can provide semantic context to a computing device when determining if an action should be performed and/or what action should be performed.

(18) Some examples provided herein relate to using a model and associated semantic embeddings to store entries in an embedding object memory that includes embeddings that correspond to content objects. A system hosting the model may be informed by semantic context and can look into the embedding object memory (e.g., a vectorized command store) to find matching content information by semantic address and/or semantic proximity (e.g., using cosine distance or another geometric n-dimensional distance function). In some examples, content objects themselves may be partitioned, sliced, and/or sub-divided by mechanisms provided herein to permit more fine-grained indexing and semantic proximity matching. This in turn can aid the discovery of useful overlap, similarity, and/or cross-connections (edges) with other content objects.

(19) The embedding object memory may store embeddings associated with models and their specific versions, which may represent the same content information in different semantic embedding spaces. When a new model is added, a content object can be re-encoded (e.g., by generating a new embedding) in the new model semantic space to add to a collection of models. In this manner, a single content object may have a locatable semantic address across models. Storing and retrieving matching content objects may require specific methodologies to ensure the content objects are available across models. The present disclosure discusses aspects of inserting entries into, retrieving information from, and rebalancing an embedding object memory.

(20) In some examples, a hierarchy may be built of the collection of models. For example, the hierarchy may be a tree, graph, or another data structure that stores content. In some examples, not only can a content object be sub-divided into more granular pieces, but sets of related content objects (such as those related by topic, time of creation, or other properties recognized by those of ordinary skill in the art) can be aggregated. The aggregated content objects can form more general higher level layers of a data structure, such as a tree. AI models, such as those described herein, can be used to create new aggregated or merged content (e.g., summary, notes, rewrites etc.) that captures higher level semantic meanings of the set below it into a single new object. The object can in turn be turned into an embedding.

(21) Some aspects of the present disclosure relate to methods, systems, and media for storing entries in an embedding object memory. Generally, one or more content items (e.g., emails, audio data, video data, messages, internet encyclopedia data, skills, commands, source code, programmatic evaluations, etc.) may be received. The one or more content items may include one or more content data (e.g., each email in the emails, each audio file in the audio data, each video file in the video data, each message in the messages, each page of the internet encyclopedia, etc.). One or more of the content data associated with the content item may be provided to one or more

semantic embedding models (e.g., a generative large language model, machine-learning model, etc.) to generate one or more semantic embeddings. One or more semantic embeddings may be received from the one or more semantic embedding models. In this respect, the large quantities of information that a computing device receives may be converted to embeddings (e.g., semantic embeddings) that can be mathematically compared and that occupy a relatively smaller amount of memory than the large quantities of information themselves.

(22) A collection of semantic embeddings may be associated with a respective semantic embedding model. For example, a first collection of embeddings may be associated with a first semantic embedding model of the one or more semantic embedding models. Further, the collection of embeddings may include a first semantic embedding generated by the first semantic embedding model for at least one content data from the respective content item. In some examples, the semantic embedding models may be optimized for specific scenarios. For example, some semantic embedding models may be configured to produce embeddings with superior results for semantic proximity on certain object types, such as emails (e.g., for a specific user and/or organization). As another example, some semantic embedding models may be configured based on specific problem types (e.g., finding documents with 3-4 word topic summaries that are relatively close based on semantic similarity). The semantic embeddings models that are optimized for specific scenarios may generate embeddings that are relatively smaller or relatively larger than other semantic embedding models that are not configured for the specific scenarios. In some examples, hints may be provided to the semantic embedding models, such as to configure the models for the specific scenarios. The hints may be provided by a user and/or by an automated system, such as to provide options or strategies related to an input provided to the models.

(23) The one or more semantic embeddings received from the semantic embedding models may be inserted into the embedding object memory. The one or more semantic embeddings may be associated with a respective indication corresponding to a location of source data associated with the one or more semantic embeddings. Further, the insertion may trigger a spatial storage operation to store a vector representation of the one or more semantic embeddings. The vector representation may be stored in at least one of an approximate nearest neighbor (ANN) tree, a k-dimensional (k-d) tree, an octree, an n-dimensional tree, or another data structure that may be recognized by those of ordinary skill in the art at least in light of teachings described herein.

(24) Additionally, or alternatively, some aspects of the present disclosure relate to methods, system, and media for retrieving information from an embedding object memory. Generally, an input embedding may be received that is generated by a machine-learning model. The input embedding is discussed in further detail later herein. A plurality of collections of stored embeddings may be retrieved by mechanisms described herein. The plurality of collections of embeddings may each correspond to respective content data. A subset of embeddings from at least one of the plurality of collections of stored embeddings may be retrieved based on a similarity to the input embedding. Further, an action may be determined based on the subset of embeddings and the input embedding.

(25) In some examples, the embedding object memory may contain tuples that include the stored embeddings, such as {value, embedding} or {key, embedding}. In such examples, the key, value, or an identification of original raw content associated with an embedding may be retrieved while the embedding stays buried within the store. Therefore, the closest embeddings to an input may be found and the associated key, value, identification, and/or reference may be returned.

(26) In some examples, original content, such as from which embeddings are generated, can be thrown away. If the embedding is generated using an auto-encoder, then given the encoding for the content (e.g., embedding), the original content can be approximately regenerated using a decoder.

(27) Some scenarios in which an embedding itself may be retrieved could include if given a set of objects, it is desirable to create a centroid (average) that represents the set of objects in average, as well as bounding coordinates that enclose a space in which the embeddings for the set of object reside. These scenarios may be useful for object aggregation and index building.

(28) Advantages of mechanisms disclosed herein may include improved accuracy for comparing subsets of large quantities of content, such as by quantitatively comparing semantic embeddings corresponding to received content. Furthermore, mechanisms disclosed herein for storing entries in an embedding object memory can improve computational efficiency by, for example, reducing an amount of memory that is needed to track content (e.g., via feature vectors or embeddings, as opposed to audio files, video files, and/or encyclopedia pages themselves). Still further, mechanisms disclosed herein can improve computational efficiency for receiving content from an embedding object memory, such as by searching for semantic embeddings that may be categorized based on their source (e.g., type of data and/or application that generated the data), as opposed to searching through relatively large quantities of raw source data stored in memory.

(29) FIG. 1 shows an example of a system **100**, in accordance with some aspects of the disclosed subject matter. The system **100** may be a system for storing entries in an embedding object memory. Additionally, or alternatively, the system **100** may be a system for using an embedding object memory, such as by retrieving information from the embedding object memory. The system **100** includes one or more computing devices **102**, one or more servers **104**, a content data source **106**, an input data source **107**, and a communication network or network **108**.

(30) The computing device **102** can receive content data **110** from the content data source **106**, which may be, for example a microphone, a camera, a global positioning system (GPS), etc. that transmits content data, a computer-executed program that generates content data, and/or memory with data stored therein corresponding to content data. The content data **110** may include visual content data, audio content data (e.g., speech or ambient noise), gaze content data, calendar entries, emails, document data (e.g., a virtual document), weather data, news data, blog data, encyclopedia data and/or other types of private and/or public content data that may be recognized by those of ordinary skill in the art. In some examples, the content data may include text, source code, commands, skills, or programmatic evaluations.

(31) The computing device **102** can further receive input data **111** from the input data source **107**, which may be, for example, a camera, a microphone, a computer-executed program that generates input data, and/or memory with data stored therein corresponding to input data. The content data **111** may be, for example, a user-input, such as a voice query, text query, etc., an image, an action performed by a user and/or a device, a computer command, a programmatic evaluation, or some other input data that may be recognized by those of ordinary skill in the art.

(32) Additionally, or alternatively, the network **108** can receive content data **110** from the content data source **106**. Additionally, or alternatively, the network **108** can receive input data **111** from the input data source **107**.

(33) Computing device **102** may include a communication system **112**, an embedding object memory insertion engine or component **114**, and/or an embedding object memory retrieval engine or component **116**. In some examples, computing device **102** can execute at least a portion of the embedding object memory insertion component **114** to generate collections of embeddings corresponding to one or more subsets of the received content data **110** to be inserted into an embedding object memory. For example, each of the subsets of the content data may be provided to a machine-learning model, such as a natural language processor and/or a visual processor, to generate a collection of embeddings. In some examples, the subsets of content data may be provided to another type of model, such as a generative large language model (LLM).

(34) Further, in some examples, computing device **102** can execute at least a portion of the embedding object memory retrieval component **116** to retrieve a plurality of collections of stored embeddings from an embedding object memory, based on an input embedding (e.g., generated based on the input data **111**), and determine an action. For example, a subset of embeddings may be retrieved from one or more collections of the plurality of collections of stored embeddings, based on a similarity to the input embedding, and the action may be determined based on the input embeddings and the subset of embeddings (e.g., as contextual information in determining an action

based on the input embeddings).

(35) Server **104** may include a communication system **118**, an embedding object memory insertion engine or component **120**, and/or an embedding object memory retrieval engine or component **122**. In some examples, server **104** can execute at least a portion of the embedding object memory insertion component **120** to generate collections of embeddings corresponding to one or more subsets of the received content data **110** to be inserted into an embedding object memory. For example, each of the subsets of the content data may be provided to a machine-learning model, such as a natural language processor and/or a visual processor, to generate a collection of embeddings. In some examples, the subsets of content data may be provided to a generative large language model (LLM).

(36) Further, in some examples, server **104** can execute at least a portion of the embedding object memory retrieval component **122** to retrieve a plurality of collections of stored embeddings from an embedding object memory, based on an input embedding (e.g., generated based on the input data **111**), and determine an action. For example, a subset of embeddings may be retrieved from one or more collections of the plurality of collections of stored embeddings, based on a similarity to the input embedding, and the action may be determined based on the input embeddings and the subset of embeddings (e.g., as contextual information in determining an action based on the input embeddings).

(37) Additionally, or alternatively, in some examples, computing device **102** can communicate data received from content data source **106** and/or input data source **107** to the server **104** over a communication network **108**, which can execute at least a portion of the embedding object memory insertion component **114** and/or the embedding object memory retrieval engine **116**. In some examples, the embedding object memory insertion component **114** may execute one or more portions of methods/processes **400** and/or **700** described below in connection with FIGS. **4** and **7**, respectively. Further in some examples, the embedding object memory retrieval component **116** may execute one or more portions of methods/processes **400** and/or **700** described below in connection with FIGS. **4** and **7**, respectively.

(38) In some examples, computing device **102** and/or server **104** can be any suitable computing device or combination of devices, such as a desktop computer, a vehicle computer, a mobile computing device (e.g., a laptop computer, a smartphone, a tablet computer, a wearable computer, etc.), a server computer, a virtual machine being executed by a physical computing device, a web server, etc. Further, in some examples, there may be a plurality of computing device **102** and/or a plurality of servers **104**. It should be recognized by those of ordinary skill in the art that content data **110** and/or input data **111** may be received at one or more of the plurality of computing devices **102** and/or one or more of the plurality of servers **104**, such that mechanisms described herein can insert entries into an embedding object memory and/or use the embedding object memory, based on an aggregation of content data **110** and/or input data **111** that is received across the computing devices **102** and/or the servers **104**.

(39) In some examples, content data source **106** can be any suitable source of content data (e.g., a microphone, a camera, a GPS, a sensor, etc.). In a more particular example, content data source **106** can include memory storing content data (e.g., local memory of computing device **102**, local memory of server **104**, cloud storage, portable memory connected to computing device **102**, portable memory connected to server **104**, etc.). In another more particular example, content data source **106** can include an application configured to generate content data. In some examples, content data source **106** can be local to computing device **102**. Additionally, or alternatively, content data source **106** can be remote from computing device **102** and can communicate content data **110** to computing device **102** (and/or server **104**) via a communication network (e.g., communication network **108**).

(40) In some examples, input data source **107** can be any suitable source of input data (e.g., a microphone, a camera, a sensor, etc.). In a more particular example, input data source **107** can

include memory storing input data (e.g., local memory of computing device **102**, local memory of server **104**, cloud storage, portable memory connected to computing device **102**, portable memory connected to server **104**, privately-accessible memory, publicly-accessible memory, etc.). In another more particular example, input data source **107** can include an application configured to generate input data. In some examples, input data source **107** can be local to computing device **102**. Additionally, or alternatively, input data source **107** can be remote from computing device **102** and can communicate input data **111** to computing device **102** (and/or server **104**) via a communication network (e.g., communication network **108**).

(41) In some examples, communication network **108** can be any suitable communication network or combination of communication networks. For example, communication network **108** can include a Wi-Fi network (which can include one or more wireless routers, one or more switches, etc.), a peer-to-peer network (e.g., a Bluetooth network), a cellular network (e.g., a 3G network, a 4G network, a 5G network, etc., complying with any suitable standard), a wired network, etc. In some examples, communication network **108** can be a local area network (LAN), a wide area network (WAN), a public network (e.g., the Internet), a private or semi-private network (e.g., a corporate or university intranet), any other suitable type of network, or any suitable combination of networks. Communication links (arrows) shown in FIG. **1** can each be any suitable communications link or combination of communication links, such as wired links, fiber optics links, Wi-Fi links, Bluetooth links, cellular links, etc.

(42) FIG. **2** illustrates examples of content, such as private content **200** and public content **250**, according to some aspects described herein. As discussed with respect to system **100**, examples described may receive content data (e.g., content data **110**) from a content data source (e.g., content data source **106**). The content data that is received may include the private content **200** and/or public content **250**. Additionally, or alternatively, the content data may include source code, commands, programmatic evaluations, or skills.

(43) Generally, when a user is interacting with a computing device (e.g., computing device **102**), they may be interacting with applications that are stored locally on the computing device and/or that can be executed locally on the computing device. Information that a user accesses or executes locally on their device may include the private content **200**.

(44) The private content includes audio content **202**, visual content **204**, gaze content **206**, calendar entries **208**, emails **210**, and documents **212**, as examples. Additional and/or alternative types of private content may be recognized by those of ordinary skill in the art.

(45) The audio content **202** may include data corresponding to speech data that is generated. For example, the audio content **202** may be generated by the computing device **102** to correspond to audio that is received from a user (e.g., where the user is speaking into a microphone a computing device that may be separate from the computing device **102**). Additionally, or alternatively, the audio content **202** may correspond to types of audio data that may be generated by a computing device, such as synthetic speech, animal sounds, beeps, buzzes, or another type of generated audio data.

(46) The visual content **204** may include data corresponding to graphical content that may be displayed and/or generated by a computing device. For example, the visual content **204** may be content that is generated via an application being run on the computing device **102** (e.g., a web-browser, a presentation application, a teleconferencing application, a business management application, etc.). The visual content **204** may include data that is scraped from a screen display of the computing device **102**. For example, any visual indication that is displayed on the computing device **102** may be included in the visual content **204**.

(47) The gaze content **206** may include data corresponding to where users are looking. For example, specific actions to be performed by a computing device may be associated with a specific location at which a user is looking and/or a combination of locations at which a user is looking within a predefined duration of time.

(48) The calendar entries **208** may include calendar data specific to one or more users. For example, the calendar data may include meetings, appointments, reservations or other types of calendar entries. Additionally, or alternatively, the calendar data may include times, locations, attendees, and/or notes regarding specific calendar entries. Additional and/or alternative data associated with calendar entries may be recognized by those of ordinary skill in the art.

(49) The emails **210** may include email data for one or more emails. For example, the emails **210** may include email data corresponding to a collection or plurality of emails. The email data may include senders and recipients, subjects, messages, images, timestamps, and/or other types of information that may be associated with emails. Additional and/or alternative data associated with calendar entries may be recognized by those of ordinary skill in the art.

(50) The virtual documents **212** may include a type of document that is found in a virtual environment. For example, the virtual document **212** may be a text-editing document, a presentation, an image, a spreadsheet, an animated series of images, a notification, or any other type of virtual document that may be recognized by those of ordinary skill in the art.

(51) In some examples, the virtual documents **212** may include interactive chats, such as chat communications between humans, between systems, and/or between humans and AI systems. Interactive chats may include dynamic streams of text that benefit from the integration of semantic memory described herein. To conduct a meaningful long running chat with the user, interaction history with an AI system (or any generated content), which is a sequence of objects, may be stored in a semantic object store, and then retrieved automatically and injected into the interactive chat's relevant history interface (e.g., a window on a computing device).

(52) Each of the plurality of types of private content **200** may be subsets of the private content **200** that may be received by mechanisms described herein, as a subset of the content data **110**. Further, while specific examples of types of private content have been discussed above, additional and/or alternative types of private content may be recognized by those of ordinary skill in the art.

(53) The public content **250** includes weather **252**, news **254**, encyclopedias **256**, blogs **258** and the like. The weather **252** may include information regarding weather that is around a user and/or at a location determined to be of interest for a user. For example, for a given time, weather information (e.g., precipitation, temperature, humidity, etc.) may be received or otherwise obtained for where a user is located (e.g., based on location content) and/or a location determined to be of interest to the user.

(54) The news **254** may include information regarding recent news stories that are determined to be of interest to a user. For example, for a given time, a relatively recent news story covering a significant event may have been released. Additional or alternative types of news stories may include holidays, birthdays, local events, national events, natural disasters, celebrity updates, scientific discoveries, sports updates, or any other type of news that may be recognized by those of ordinary skill in the art.

(55) The encyclopedia **256** may include publicly available encyclopedia information. For example, the encyclopedia **256** may include information from an online database of encyclopedic information. Additionally, or alternatively, the encyclopedia **256** may include pages from an online encyclopedia website. Additional or alternative types of encyclopedia information may be recognized by those of ordinary skill in the art.

(56) The blogs **256** may include information from blogs. For example, the blogs may include publicly available posts from users of a blog website and/or a social media platform. The blogs may be posted by, for example, famous people, such as chefs, politicians, actors, etc. Alternatively, the blogs may be posted by other users who post content online that may be publicly accessible by mechanisms disclosed herein.

(57) Generally, the different content types discussed with respect to FIG. 2 provide various types of content that may be received or otherwise accessed by a computing device and that may be useful in providing contextual information for determining actions described herein. Further, while

specific subsets of content were described above with respect to one of the private content **200** and the public content **250**, it should be recognized that in some examples the subsets of content may instead be described with respect to the other of the private content **200** or the public content **250**. Further, it is noted that additional and/or alternative types of private content **200** and/or public content **250** will be recognized by those of ordinary skill in the art.

(58) FIG. 3 illustrates an example flow **300** of storing entries in an embedding object memory. The example flow begins with a first set of content data **302** being received. The first set of content **302** may be received from a content data source, such as the content data source **110** described earlier herein with respect to FIG. 1. Further, the first set of content **302** may be one of private content or public content, such the private content **200** and/or the public content **250** described earlier herein with respect to FIG. 2. The first set of content **302** may include one or more content items or objects that include content data.

(59) The first set of content data **302** may be input into a first model **304**. For example, the first model **304** may include a machine-learning model, such as a machine-learning model that is trained to generate one or more embeddings, based on received content data. In some examples, the first model **304** includes a generative model, such as a generative large language model (LLM). In some examples, the first model **304** includes a natural language processor. Alternatively, in some examples, the first model **304** includes a visual processor. The first model **304** may be trained on one or more datasets that are compiled by individuals and/or systems. Additionally, or alternatively, the first model **304** may be trained based on a dataset that includes information obtained from the Internet. Further, the first model **304** may include a version. One of skill in the art will appreciate that any type of model may be employed as part of the aspects disclosed herein without departing from the scope of this disclosure.

(60) The first model **304** may output a first collection of embeddings **306**. The first collection of embeddings **306** may include one or more semantic embeddings. The first collection of embeddings **306** may be unique to the first model **304**. Additionally, or alternatively, the first collection of embeddings **306** may be unique to the specific version of the first model **304**. Each embedding in the first collection of embeddings **306** may correspond to a respective content object and/or respective content data from the first set of content data **302**. Therefore, for each object and/or content data in the first set of content data **302**, the first model **304** may generate an embedding.

(61) Each of the embeddings in the first collection of embeddings **306** may be associated with a respective indication (e.g., byte address) corresponding to a location of source data associated with the embeddings. For example, if the first set of content data **302** corresponds to a set of emails, then the first collection of embeddings **306** may be hash values generated for each of the emails based on the content within the emails (e.g., abstract meaning of words or images included in the email). An indication corresponding to the location of an email (e.g., in memory) may then be associated with the embedding generated based off of the email. In this respect, the actual source data of a content object may be stored separately from the corresponding embedding that may occupy less memory.

(62) The first collection of embeddings **306** and indications of corresponding source data associated with the first set of content data **302** (and thereby associated with the first collection of embeddings **306**) may be stored together, such as in a first storage object **308**. In some examples, the embeddings **306** and indications of corresponding source data may be stored in an embedding object memory **310**. In some examples, the embeddings **306** may be a list of floating numbers, but the embeddings **306** may be linked to original source data associated with the embeddings **306** (e.g., the first set of content data **302**), such that mechanisms provided herein can locate, retrieve, and/or otherwise access source data (e.g., an email thread, a text document, etc.) based on one or more embeddings from the embeddings **306**.

(63) The example flow **300** further includes with a second set of content data **303** being received.

The second set of content data **303** may be received from a different content data source than the first set of content data **302**. For example, the first set of content data **302** may be audio data and the second set of content data **303** may be video data. As another example, the first set of content data **302** may correspond to a first application (e.g., a word-processing application, a spreadsheet application, etc.) and the second set of content data **303** may correspond to a second application that is different than the first application. Additional and/or alternative combinations and examples of content data will be recognized by those of ordinary skill in the art.

(64) The second set of content data **303** may be input into a second model **305**. For example, the second model **305** may be a machine-learning model, such as a machine-learning model that is trained to generate one or more embeddings, based on received content data. In some examples, the second model **305** includes a generative model, such as a generative large language model (LLM). In some examples, the second model **305** includes a natural language processor. Alternatively, in some examples, the second model **305** includes a visual processor. The model may be trained on one or more datasets that are compiled by individuals and/or systems. Additionally, or alternatively, the second model **305** may be trained based on a dataset that includes information obtained from the Internet. Further, the second model **305** may include a version. One of skill in the art will appreciate that any type of model may be employed as part of the aspects disclosed herein without departing from the scope of this disclosure.

(65) The second model **305** may output a second collection of embeddings **307**. The second collection of embeddings **307** may include one or more semantic embeddings. The second collection of embeddings **307** may be unique to the second model **305**. Additionally, or alternatively, the second collection of embeddings **307** may be unique to the specific version of the second model **305**. For example, if the first model **304** and the second model **305** were both to receive the same set of content data, they may output different embeddings. Further, if a first version of the second model **305** received the same set of content data as a second version of the second model **305**, then the first and second version of the second model **305** may output different collections of embeddings. Each embedding in the second collection of embeddings **307** may correspond to a respective content object and/or respective content data from the second set of content data **303**. Therefore, for each object and/or content data in the second set of content data **303**, the second model **305** may generate a respective embedding.

(66) Each of the embeddings in the second collection of embeddings **307** may be associated with a respective indication (e.g., byte address, key, file name) corresponding to a location of source data associated with the embeddings. For example, if the second set of content data **303** corresponds to a set of encyclopedia data, then the second collection of embeddings **307** may be hash values generated for each of the pages of the encyclopedia based on the content within the pages (e.g., abstract meaning of words or images included in the pages). An indication corresponding to the location of the encyclopedia pages may then be associated with the embedding generated based off of the encyclopedia data. In this respect, the actual source data of a content object may be stored separately from the corresponding embedding that may occupy less memory.

(67) The second collection of embeddings **307** and indications of corresponding source data associated with the second set of content data **303** (and thereby associated with the second collection of embeddings **307**) may be stored together, such as in a second storage object **309**. In some examples, the embeddings **307** and indications of corresponding source data may be stored in an embedding object memory **310**. In some examples, the embeddings **307** may be a list of floating numbers, but the embeddings **307** may be linked to original source data associated with the embeddings **307** (e.g., the second set of content data **303**), such that mechanisms provided herein can locate, retrieve, and/or otherwise access source data (e.g., an email thread, a text document, etc.) based on one or more embeddings from the embeddings **307**.

(68) Accordingly, the embedding object memory **310** may store a plurality of collections of embeddings (e.g., the first collection of embeddings **306** and the second collection of embeddings

307) with indications that correspond to locations of source data associated with the one or more embeddings. The plurality of collections of embeddings may be a plurality of semantic embeddings. The indications may be byte addresses, uniform resource indicators (e.g., uniform resource links), or another form of data that is capable of identifying a location of source data. The source data may be stored in memory of a computing device or server on which the embedding object memory **310** is located. Additionally, or alternatively, the source data may be stored in memory that is remote from a computing device or server on which the embedding object memory **310** is located.

(69) FIG. **4** illustrates an example method **400** for storing an entry in an embedding object memory, according to some aspects described herein. In examples, aspects of method **400** are performed by a device, such as computing device **102** and/or server **104**, discussed above with respect to FIG. **1**.

(70) Method **400** begins at operation **402** wherein one or more content items objects are received. The content items may have one or more content data. The one or more content data may include at least one public content data and/or at least one private content data. The content data may be similar to the content data **110** discussed with respect to FIG. **1**. Additionally, or alternatively, the content data may be similar to the private content data **200** and/or public content data **250** discussed with respect to FIG. **2**. In some examples, the content data may include at least one of a skill, command, or programmatic evaluation.

(71) The content data may include one or more of audio content data, visual content data, gaze content data, calendar entries, emails, virtual documents, data generated by specific software applications, weather content data, news content data, encyclopedia content data, and/or blog content data. Additional and/or alternative types of content data may be recognized by those of ordinary skill in the art.

(72) At operation **404**, it is determined if at least one of the subsets of content data has an associated semantic embedding model. For example, the semantic embedding model may be similar to the first model **304** and/or the second model **305** discussed with respect to FIG. **3**. The semantic embedding model may be trained to generate one or more embeddings, based on the at least one of the subsets of content data. In some examples, the semantic embedding model may include a natural language processor. In some examples, the semantic embedding model may include a visual processor. In some examples, the semantic embedding model may include a machine-learning model. Still further, in some examples, the semantic embedding model may include a generative large language model. The semantic embedding model may be trained on one or more datasets that are compiled by individuals and/or systems. Additionally, or alternatively, the semantic embedding model may be trained based on a dataset that includes information obtained from the Internet.

(73) If it is determined that each of the content data do not have an associated semantic embedding model, flow branches “NO” to operation **406**, where a default action is performed. For example, the content data and/or content items may have an associated pre-configured action. In other examples, method **400** may comprise determining whether the content data and/or content items have an associated default action, such that, in some instances, no action may be performed as a result of the received content items. Method **400** may terminate at operation **406**. Alternatively, method **400** may return to operation **402** to provide an iterative loop of receiving one or more content items and determining if at least one of the content data of the content items have an associated semantic embedding model.

(74) If however, it is determined that at least one of the content data have an semantic embedding model, flow instead branches “YES” to operation **408**, where, one of the content data associated with the content item are provided to one or more semantic embedding models. The semantic embedding models generate one or more semantic embeddings. Further, the one or more semantic embedding models may include a version. Each of the semantic embeddings generated by the respective semantic embedding model may include metadata corresponding to the version of the

semantic embedding model that generated the embedding. The embedding object memory may be associated with a versioned model store (or a service that stores the model and version, and can run inference processes based thereon).

(75) In some examples, the content data is provided to the one or more semantic embedding models locally. In some examples, the content data is provided to the one or more semantic embedding model via an application programming interface (API). For example, a first device (e.g., computing device **102** and/or server **104**) may interface with a second device (e.g., computing device **102** and/or server **104**) via an API that is configured to provide embeddings or indications thereof, in response to receiving content data or indications thereof.

(76) At operation **410**, one or more semantic embeddings are received from one or more of the semantic embedding models. In some examples, the one or more semantic embeddings are a collection or plurality of embeddings. For example, a collection of embeddings (such as the first collection of embeddings **306**) may be associated with a first semantic embedding model (e.g., the first model **304**). In some examples, the collection of embeddings may be uniquely associated with the first semantic embedding model, such that a second semantic embedding model generates a different collection of embeddings than the first semantic embedding model. Further, in some examples, the collection of embeddings may be uniquely associated with the version of the first semantic embedding model, such that a different version of the first semantic embedding model generates a different collection of embeddings.

(77) The collection of embeddings may include an embedding generated by the first semantic embedding model for at least one content data of a plurality of content data. For example, the content data may correspond to an email, an audio file, a message, a website page, etc. Additional and/or alternative types of content objects or items that correspond to the content data may be recognized by those of ordinary skill in the art.

(78) At operation **412**, the one or more semantic embeddings are inserted into the embedding object memory. In some examples, the semantic object memory may be an index, or a database, or a tree (e.g., ANN tree, k-d tree, etc.), or another type of memory that may be recognized by those of ordinary skill in the art. The embedding object memory includes one or more embeddings (e.g., semantic embeddings) from the collection of embeddings. Further, the one or more embeddings are associated with a respective indication that corresponds to a location of source data associated with the one or more embeddings. The source data can include one or more of audio files, text files, image files, video files, and/or website pages. Additional and/or alternative types of source data may be recognized by those of ordinary skill in the art. The embeddings may occupy relatively little memory. For example, each embedding may be a 64-bit hash. Comparatively, the source data may occupy a relatively large amount of memory. Accordingly, by storing an indication to the location of the source data in the embedding object memory, as opposed to the raw source data, mechanisms described herein may be relatively efficient for memory usage of one or more computing devices. In some examples, ordinary or cryptographic hashes may not be semantically useful for techniques provided herein because they do not support proximity search, and content may be distributed randomly all over a 64-bit space. However, in some examples, locality sensitive hashing (LSH) would be semantically useful.

(79) The indications that corresponds to a location of source data may be byte addresses, uniform resource indicators (e.g., uniform resource links), or another form of data that is capable of identifying a location of source data. Further, the embedding object memory may be stored at a location that is different than the location of the source data. Additionally, or alternatively, the source data may be stored in memory of a computing device or server on which the embedding object memory (e.g., embedding object memory **310**) is located. Additionally, or alternatively, the source data may be stored in memory that is remote from a computing device or server on which the embedding object memory is located.

(80) In some examples, the insertion at operation **412** triggers a spatial storage operation to store a

vector representation of the one or more embeddings (e.g., the one or more semantic embeddings). The vector representation may be stored in a data structure, such as an ANN tree, a k-d tree, an n-dimensional tree, an octree, or another data structure that may be recognized by those of ordinary skill in the art in light of teachings described herein. Additional and/or alternative types of storage mechanisms that are capable of storing vector space representations may be recognized by those of ordinary skill in the art.

(81) At operation **414**, the embedding object memory is provided. For example, the embedding object memory may be provided as an output for further processing to occur. In some examples, the embedding object memory may be used to retrieve information and/or generate an action, as will be described in some examples in further detail herein.

(82) In some examples, the method **400** further includes providing an updated semantic embedding model to replace at least one of the semantic embedding models. The updated semantic embedding model includes an updated version that is different than the version of the at least one of the semantic embedding models. From the updated semantic embedding model, an updated collection of embeddings may be received that correspond to the collections of embeddings that are generated by the at least one of the semantic embedding models. Further, the updated collection of embeddings may be stored in the embedding object memory with metadata that corresponds to the updated version.

(83) In some examples a method for updating the embedding object memory based on adding or removing a model may be provided that is similar in concept to adding a new index to a database or a new search index to a data source. An example of such a method may include retrieving content data associated with a first collection of embeddings generated by a first model, inputting the retrieved content data into a second model (e.g., which is a new version of the first model) to generate a second collection of embeddings, wherein the first and second collection of embeddings are both retrievable with indications corresponding to the model by which they were generated. In some examples, mechanisms provided herein can crawl an original set of content data (a reference to which could be retained) and create a new embedding collection. In some examples, mechanisms provided herein can retrieve keys associated with each embedding, then retrieve the content data for each key, and create a new embedding. Such techniques can be helpful for object stores (e.g., a semantic object memory) with Get/Put semantics. Updating the embedding object memory may be easiest if raw content items are stored co-located with embeddings and/or embedding indexes.

(84) Generally, embeddings may be used as quantitative representations of abstract meaning that is discerned from content data. Accordingly, different semantic embedding models may generate different embeddings for the same received content data, such as depending on how the semantic embedding models are trained (e.g., based on different data sets and/or training methods) or otherwise configured. In some examples, some semantic embedding models may be trained to provide relatively broader interpretations of received content than other semantic embedding models that are trained to provide relatively narrower interpretations of content data. The ability to provide relatively broader or narrower interpretations can also be accomplished via aggregation of the source objects (discussed earlier herein). Accordingly, such semantic embedding models may generate different embeddings.

(85) In some examples, the method **400** further includes deleting the collection of embeddings that correspond to the updated collection of embeddings. The collection of embeddings include metadata that corresponds to a version that is different than the updated version. Generally, if a first version of a semantic embedding model is replaced with a second version of a semantic embedding model, then it may be efficient for memory storage to delete embeddings generated by the first version of the semantic embedding model since new embeddings may be stored that are generated by the second version of the semantic embedding model. For example, the embeddings generated by the second version of the semantic embedding model (e.g., a newer version) may be more desirable for use in further processing than the embeddings generated by the first version of the

semantic embedding model (e.g., an older version).

(86) One of ordinary skill in the art will recognize that in some examples embeddings may not be deleted, such as when it is not necessary or desirable to delete embeddings that correspond to different versions of a semantic embedding model. Further, in some examples, embeddings that correspond to different versions of a semantic embedding model may be stored in the embedding object memory with options to roll-back and/or roll-forward to different version of embeddings. For example, if a version of embeddings is corrupted or unusable, then it may be desirable to conduct further processing using a different version of embeddings that is accessible and/or useable.

(87) Method **400** may terminate at operation **414**. Alternatively, method **400** may return to operation **402** (or any other operation from method **400**) to provide an iterative loop, such as of receiving one or more content items and inserting one or more semantic embeddings into an embedding object memory based thereon.

(88) FIG. 5 illustrates an example flow **500** for retrieving information from an embedding object memory, according to some aspects described herein. The example flow begins with receiving an input embedding **502**. The input imbedding **502** may be generated by a model, such as a machine-learning model. For example, user-input may be received by a system and the user-input may be input into an embedding generator (e.g., a model, which may be similar to the first model **304** and/or the second model **305** described with respect to FIG. 3). The input embedding **502** may correspond to an intent of a user. For example, a user may intend to perform an action with their computing device, such as book a flight, or send an email, or create a presentation, or locate a file, or run a program. Additional and/or alternative intents that a user has with respect to a computing device may be recognized by those of ordinary skill in the art.

(89) In some examples, the input embedding **502** may be generated based on something other than user-input. For example, the input embedding **502** may be generated based on an algorithm that triggers the input embedding **502** to be generated. Additionally, or alternatively, the input embedding **502** may be generated based on an audio and/or visual input that is received by a computer (e.g., independent from user-input). In some example, the input embedding **502** may be generated by a natural language processor, or a visual processor, or a generative large language model. Additional and/or alternative methods for generating the input embedding **502** may be recognized by those of ordinary skill in the art.

(90) The input embedding **502** may be input into a semantic retrieval engine **504**. The semantic retrieval engine or component **504** may include hardware and/or software that retrieves, from a embedding object memory **506** (e.g., which may be similar to the embedding object memory provided by method **400**), a plurality of collections of stored embeddings, such as a first collection of stored embeddings **508**, a second collection of stored embeddings **510**, and a third collection of stored embeddings **512**. The first, second, and third collections of stored embeddings **508**, **510**, and **512** may each be retrieved based on the input embedding **502**.

(91) Each of the collections of stored embeddings **508**, **510**, **512** may correspond to respective content data, such as content data **110**, private content data **200**, and/or public content data **250**. For example, the first collection of stored embeddings **508** may include embeddings (e.g., semantic embeddings) that correspond to calendar invites, the second collection of stored embeddings **510** may include embeddings (e.g., semantic embeddings) that correspond to emails, and the third collection of stored embeddings **512** may include embeddings (e.g., semantic embeddings) that correspond to audio files. Additional and/or alternative content types to which the collections of stored embeddings **508**, **510**, and **512** correspond may be recognized by those of ordinary skill in the art.

(92) A subset of embeddings may be retrieved from at least one of the plurality of collections of stored embeddings **508**, **510**, **512** based on a similarity to the input embedding **502**. For example, embeddings that correspond to one or more emails, one or more calendar invites, one or more audio

files, one or more video files, etc. may be received from respective collections of stored embeddings, such as the plurality of collections of stored embeddings **508, 510, 512**.

(93) In some examples, a respective similarity between the input embedding and each embedding of the plurality of collections of stored embeddings may be determined. The determined similarities may be compared to a predetermined threshold. Further, the subset of embeddings with similarities to the input embedding that are less than the predetermined threshold may be retrieved. Therefore, subsets of embeddings from at least one of the collections of embeddings that are determined to be related to the input embedding may be retrieved. Mechanisms for determining a similarity of embeddings to the input embeddings are discussed in further detail with respect to the vector space **600** of FIG. **6**.

(94) In some examples, an ordered ranking of the similarities may be determined. For example, the similarities may be ordered (e.g., in an increasing or decreasing fashion) based on values associated with the determined similarities. Therefore, a subset of embeddings may be retrieved based on the ordered ranking (e.g., a specified number of top and/or bottom rankings may be retrieved). In some examples, a subset of semantic embeddings may be retrieved with similarities to the input embedding that are less than the predetermined threshold and/or based on the ordered ranking.

(95) At action determination engine or component **514**, a semantic-based action or action **516** is determined based on the subsets of embeddings (e.g., from the plurality of collections of stored embeddings **508, 510, 512**) and the input embedding **502**. The action **516** may further be determined based on source data **518**. For example, mechanisms described herein may locate raw source data corresponding to the subset of embeddings (e.g., similar as to how, in FIG. **3**, the first set of content data **302** may include raw source data associated with the embeddings **306**).

Mechanisms may further determine the action based on further analysis of the content of the raw source data. As discussed earlier herein, the source data may include one or more of audio files, text files, image files, video files, website pages, blog posts, data generated by specific software applications, etc.

(96) The semantic-based action may be any of a plurality of different actions. For example, the action may be an action that is performed by a user and/or by a system. The action may include instructions that are output to a user based on the input embedding **502** and the embedding object memory **506**. Additionally or alternatively, the action may be performed automatically by a computing device (e.g., computing device **102**). Some examples of actions include generating a draft email, sending a calendar invite, running a simulation, executing a specific set of operations via at least one processor, scheduling a meeting, booking a flight, leaving a voicemail, etc.

(97) In one particular example, the input embedding **502** may be generated based on a user-input indicating "I want to book a flight." Based on the input embedding **502**, mechanisms described herein may retrieve a plurality of collections of stored embeddings that correspond to emails, audio message, and calendar invites. A subset of the email embeddings may correspond to email source data content such as "I look forward to meeting you in New York next Thursday." A subset of audio messages may correspond to audio source data content such as "Let's plan something when you get back from New York on Monday." A subset of calendar invites may correspond to calendar invite source data such as "Trip to the NYC office from Thursday to Monday." Accordingly, based on the subset of embeddings, the intent embedding, and/or the source data, the semantic-based action may include automatically booking a flight to New York (e.g., departing Thursday and returning on Monday). Additionally, or alternatively, the action may include outputting a list of potential flights that a user may desire to book.

(98) In another particular example, the input embedding **502** may be generated based on a system reading an email stating, "Please prepare a market analysis presentation for Client A." Based on the input embedding **502**, mechanisms described herein may retrieve a plurality of collections of stored embeddings that correspond to web pages, news sources, emails, etc. A subset of the web page embeddings may correspond to source data that includes publicly available web pages about Client

A and its competitors. A subset of the news source embeddings may correspond to source data that includes recent articles about Client A and/or its competitors. A subset of the email embeddings may correspond to source data that includes emails about the market sector in which Client A operates. Accordingly, based on the subset of embeddings, the intent embedding, and/or the source data, the semantic-based action may include generating a market analysis presentation. Additionally, or alternatively, the action may include generating graphics for a market analysis presentation. Additionally, or alternatively, the action may include outputting information that a user may rely on for creating a market analysis presentation.

(99) While particular examples of actions have been provided above, the particular examples are merely examples. Additional and/or alternative examples should be recognized by those of ordinary skill in the art in light, at least in light of teachings provided herein.

(100) FIG. 6 illustrates an example vector space **600** according to some aspects described herein. The vector space **600** includes a plurality of feature vectors, such as a first feature vector **602**, a second feature vector **604**, a third feature vector **606**, a fourth feature vector **608**, and a fifth feature vector **610**. Each of the plurality of feature vectors **602**, **604**, **606**, and **608** correspond to a respective embedding **603**, **605**, **607**, **609** generated based on a plurality of subsets of content data (e.g., subsets of content data **110**, private content **200**, and/or public content **250**). The embeddings **603**, **605**, **607**, and **609** may be semantic embeddings. The fifth feature vector **610** is generated based on an input embedding **611**. The embeddings **603**, **605**, **607**, **609** may be similar to the first collection of embeddings **306** and/or the second collection of embeddings **307**, discussed with respect to FIG. 3. The input embedding **611** may be similar to the input embedding **502**.

(101) The feature vectors **602**, **604**, **606**, **608**, **610** each have distances that are measurable between each other. For example, a distance between the feature vectors **602**, **604**, **606**, and **608** and the fifth feature vector **610** corresponding to the input embedding **611** may be measured using cosine similarity. Alternatively, a distance between the feature vectors **602**, **604**, **606**, **608** and the fifth feature vector **610** may be measured using another distance measuring technique (e.g., an n-dimensional distance function) that may be recognized by those of ordinary skill in the art.

(102) A similarity of each of the feature vectors **602**, **604**, **606**, **608** to the feature vector **610** corresponding to the input embedding **611** may be determined, for example based on the measured distances between the feature vectors **602**, **604**, **606**, **608** and the feature vector **610**. The similarity between the feature vectors **602**, **604**, **606**, **608** and the feature vector **610** may be used to group or cluster the feature vectors **602**, **604**, **606**, and **608** in one or more collections of feature vectors, such as a collection **612**, thereby generating a collection of embeddings.

(103) In some examples, the collection **612** may include a predetermined number of feature vectors, such that groups of feature vectors are given a predetermined size. Additionally, or alternatively, in some examples, The distances between each of the feature vectors **602**, **604**, **606**, **608** and the feature vector **610** corresponding to the input embedding **611** may be compared to a predetermined threshold.

(104) The embeddings **603** and **605** that correspond to feature vectors **602** and **604**, respectively, may fall within the same content group. For example, the embedding **603** may be related to a first email, and the embedding **605** may be related to a second email. Additional and/or alternative examples of content groups in which the embeddings may be categorized may be recognized by those of ordinary skill in the art.

(105) The collection **612** may be stored in a data structure, such as an ANN tree, a k-d tree, an octree, another n-dimensional tree, or another data structure that may be recognized by those of ordinary skill in the art that is capable of storing vector space representations. Further, memory corresponding to the data structure in which the collection **612** is stored may be arranged or stored in a manner that groups the embeddings and/or vectors in the collection **612** together, within the data structure. In some examples, feature vectors and their corresponding embeddings generated in accordance with mechanisms described herein may be stored for an indefinite period of time.

Additionally, or alternatively, in some examples, as new feature vectors and/or embeddings are generated and stored, the new feature vectors and/or embeddings may overwrite older feature vectors and/or embeddings that are stored in memory (e.g., based on metadata of the embeddings indicating a version), such as to improve memory capacity. Additionally, or alternatively, in some examples, feature vectors and/or embeddings may be deleted from memory at specified intervals of time, and/or based on an amount of memory that is available (e.g., in the embedding object memory **310**), to improve memory capacity.

(106) Generally, the ability to store embeddings corresponding to received content data allows a user to associate and locate data in a novel manner that has the benefit of being computationally efficient. For example, instead of storing a video recording of a screen of a computing device, or web pages on the Internet, a user may instead store, using mechanisms described herein, embeddings corresponding to content objects. The embeddings may be hashes, as opposed to, for example, video recordings that may be a few hundreds of thousands of pixels, per frame. Therefore, the mechanisms described herein are efficiency for reducing memory usage, as well as for reducing usage of processing resources to search through stored content. Additional and/or alternative advantages may be recognized by those of ordinary skill in the art.

(107) FIG. 7 illustrates an example method **700** for retrieving information from an embedding object memory, according to some aspects described herein. In examples, aspects of method **700** are performed by a device, such as computing device **102** and/or server **104**, discussed above with respect to FIG. 1.

(108) Method **700** begins at operation **702** wherein an input embedding is received. The input imbedding may be generated by model, such as a machine-learning model. For example, user-input may be received and the input embedding may be generated based on the user-input. The input embedding **502** may correspond to an intent of a user. For example, a user may intend to perform an action with their computing device, such as book a flight, or send an email, or create a presentation, or locate a file, or run a program. Additional and/or alternative intents that a user has with respect to a computing device may be recognized by those of ordinary skill in the art.

(109) In some examples, the input embedding may be generated based on something other than user-input. For example, the input embedding may be generated based on an algorithm that triggers the input embedding to be generated. Additionally, or alternatively, the input embedding may be generated based on an audio and/or visual input that is received by a computer (e.g., independent from user-input). Additional and/or alternative methods for generating the input embedding, such as based on received and/or generated data, may be recognized by those of ordinary skill in the art.

(110) At operation **704**, a plurality of collections of stored embeddings (e.g., semantic embeddings) are retrieved. The plurality of collections of stored embeddings may be retrieved from an embedding object memory (e.g., the embedding object memory provided by method **400** illustrated in FIG. 4). The plurality of collections of stored embeddings may be retrieved based on and/or in response to the input embedding. Each of the plurality of collections of stored embeddings may correspond to respective content data, such as content data **110**, private content data **200**, and/or public content data **250**. For example, the collections of stored embeddings may include embeddings that correspond to calendar invites, emails, audio files, and/or data generated by specific software applications. Additional and/or alternative content types to which the collections of stored embeddings correspond may be recognized by those of ordinary skill in the art.

(111) At operation **706**, a subset of embeddings (e.g., subset of semantic embeddings) from at least one of the plurality of collections of stored embeddings are retrieved. For example, the subset of embeddings may be retrieved based on a similarity to the input embedding. A respective similarity between the input embedding and each embedding of the plurality of collections of stored embeddings may be determined. For example, the similarity may be based on a distance between the embeddings in a vector space, such as the vector space **600** described with respect to FIG. 6. Each of the similarities may be compared to a predetermined threshold, such that if a subset of

embeddings have similarities that are greater than and/or less than the predetermined threshold, the subset of embeddings may be retrieved. Therefore, the subsets of embeddings from at least one of the collections of embeddings that are determined to be related (e.g., significantly related, based on a comparison to the predetermined threshold) to the input embedding may be retrieved.

(112) At operation **708**, it is determined if there is an action associated with the subset of embeddings and the input embedding. For example, the action may be determined based on the subset of imbedding and the input embedding. In some examples, source data that is associated with the subset of embeddings may be located (e.g., local to a device on which method **700** is being executed and/or remote from a device on which method **700** is being executed) and the action may be further determined based on the source data. As discussed earlier herein, the source data may include one or more of audio files, text files, image files, video files, website pages, blog posts, news stories, data generated by specific software applications, etc.

(113) If it is determined that there is not an action associated with the subset of embeddings, flow branches “NO” to operation **710**, where a default action is performed. For example, the subset of embeddings may have an associated pre-configured action. In other examples, method **700** may comprise determining whether the subset of embeddings has an associated default action, such that, in some instances, no action may be performed as a result of the retrieved subset of embeddings. Method **700** may terminate at operation **710**. Alternatively, method **700** may return to operation **702** to provide an iterative loop of receiving an input embedding, retrieving a plurality of collections of stored embeddings, retrieving a subset of embeddings from at least one of the plurality of collections of stored embeddings, and determining if there is an action associated with the subset of embeddings and the input embedding.

(114) If however, it is determined that there is an action associated with the subset of embeddings and the input embedding, flow instead branches “YES” to operation **712**, where the action is provided as an output. For example, the action may be provided as an output to a user, a system on which method **700** is being executed, and/or a system remote from that on which method **700** is being executed. Further, in some examples, the method **700** may further include adapting a computing device to perform the action that is provided.

(115) The action may be any of a plurality of different actions. For example, the action may be an action that is performed by a user and/or by a system. The action may include instructions and/or information that are output to a user. Additionally or alternatively, the action may be performed automatically by a computing device (e.g., computing device **102**). Some examples of actions include generating a draft email, sending a calendar invite, running a simulation, executing a specific set of operations via at least one processor, scheduling a meeting, booking a flight, leaving a voicemail, etc. While particular examples of actions have been provided herein, the particular examples are merely examples. Additional and/or alternative examples should be recognized by those of ordinary skill in the art, at least in light of teachings provided herein.

(116) FIGS. **8A** and **8B** illustrate overviews of an example generative machine learning model that may be used according to aspects described herein. With reference first to FIG. **8A**, conceptual diagram **800** depicts an overview of pre-trained generative model package **804** that processes an input **802** to generate model output for storing entries in and/or retrieving information from an embedding object memory **806** according to aspects described herein. Examples of pre-trained generative model package **804** includes, but is not limited to, Megatron-Turing Natural Language Generation model (MT-NLG), Generative Pre-trained Transformer 3 (GPT-3), Generative Pre-trained Transformer 4 (GPT-4), BigScience BLOOM (Large Open-science Open-access Multilingual Language Model), DALL-E, DALL-E 2, Stable Diffusion, or Jukebox.

(117) In examples, generative model package **804** is pre-trained according to a variety of inputs (e.g., a variety of human languages, a variety of programming languages, and/or a variety of content types) and therefore need not be finetuned or trained for a specific scenario. Rather, generative model package **804** may be more generally pre-trained, such that input **802** includes a

prompt that is generated, selected, or otherwise engineered to induce generative model package **804** to produce certain generative model output **806**. For example, a prompt includes a context and/or one or more completion prefixes that thus preload generative model package **804** accordingly. As a result, generative model package **804** is induced to generate output based on the prompt that includes a predicted sequence of tokens (e.g., up to a token limit of generative model package **804**) relating to the prompt. In examples, the predicted sequence of tokens is further processed (e.g., by output decoding **816**) to yield output **806**. For instance, each token is processed to identify a corresponding word, word fragment, or other content that forms at least a part of output **806**. It will be appreciated that input **802** and generative model output **806** may each include any of a variety of content types, including, but not limited to, text output, image output, audio output, video output, programmatic output, and/or binary output, among other examples. In examples, input **802** and generative model output **806** may have different content types, as may be the case when generative model package **804** includes a generative multimodal machine learning model.

(118) As such, generative model package **804** may be used in any of a variety of scenarios and, further, a different generative model package may be used in place of generative model package **804** without substantially modifying other associated aspects (e.g., similar to those described herein with respect to FIGS. 1-7). Accordingly, generative model package **804** operates as a tool with which machine learning processing is performed, in which certain inputs **802** to generative model package **804** are programmatically generated or otherwise determined, thereby causing generative model package **804** to produce model output **806** that may subsequently be used for further processing.

(119) Generative model package **804** may be provided or otherwise used according to any of a variety of paradigms. For example, generative model package **804** may be used local to a computing device (e.g., computing device **102** in FIG. 1) or may be accessed remotely from a machine learning service. In other examples, aspects of generative model package **804** are distributed across multiple computing devices. In some instances, generative model package **804** is accessible via an application programming interface (API), as may be provided by an operating system of the computing device and/or by the machine learning service, among other examples.

(120) With reference now to the illustrated aspects of generative model package **804**, generative model package **804** includes input tokenization **808**, input embedding **810**, model layers **812**, output layer **814**, and output decoding **816**. In examples, input tokenization **808** processes input **802** to generate input embedding **810**, which includes a sequence of symbol representations that corresponds to input **802**. Accordingly, input embedding **810** is processed by model layers **812**, output layer **814**, and output decoding **816** to produce model output **806**. An example architecture corresponding to generative model package **804** is depicted in FIG. 8B, which is discussed below in further detail. Even so, it will be appreciated that the architectures that are illustrated and described herein are not to be taken in a limiting sense and, in other examples, any of a variety of other architectures may be used.

(121) FIG. 8B is a conceptual diagram that depicts an example architecture **850** of a pre-trained generative machine learning model that may be used according to aspects described herein. As noted above, any of a variety of alternative architectures and corresponding ML models may be used in other examples without departing from the aspects described herein.

(122) As illustrated, architecture **850** processes input **802** to produce generative model output **806**, aspects of which were discussed above with respect to FIG. 8A. Architecture **850** is depicted as a transformer model that includes encoder **852** and decoder **854**. Encoder **852** processes input embedding **858** (aspects of which may be similar to input embedding **810** in FIG. 8A), which includes a sequence of symbol representations that corresponds to input **856**. In examples, input **856** includes input content **802** corresponding to a type of content, aspects of which may be similar to input data **111**, private content **200**, public content **250**, first set of content data **302**, and/or second set of content data **303**.

(123) Further, positional encoding **860** may introduce information about the relative and/or absolute position for tokens of input embedding **858**. Similarly, output embedding **874** includes a sequence of symbol representations that correspond to output **872**, while positional encoding **876** may similarly introduce information about the relative and/or absolute position for tokens of output embedding **874**.

(124) As illustrated, encoder **852** includes example layer **870**. It will be appreciated that any number of such layers may be used, and that the depicted architecture is simplified for illustrative purposes. Example layer **870** includes two sub-layers: multi-head attention layer **862** and feed forward layer **866**. In examples, a residual connection is included around each layer **862**, **866**, after which normalization layers **864** and **868**, respectively, are included. Decoder **854** includes example layer **890**. Similar to encoder **852**, any number of such layers may be used in other examples, and the depicted architecture of decoder **854** is simplified for illustrative purposes. As illustrated, example layer **890** includes three sub-layers: masked multi-head attention layer **878**, multi-head attention layer **882**, and feed forward layer **886**. Aspects of multi-head attention layer **882** and feed forward layer **886** may be similar to those discussed above with respect to multi-head attention layer **862** and feed forward layer **866**, respectively. Additionally, masked multi-head attention layer **878** performs multi-head attention over the output of encoder **852** (e.g., output **872**). In examples, masked multi-head attention layer **878** prevents positions from attending to subsequent positions. Such masking, combined with offsetting the embeddings (e.g., by one position, as illustrated by multi-head attention layer **882**), may ensure that a prediction for a given position depends on known output for one or more positions that are less than the given position. As illustrated, residual connections are also included around layers **878**, **882**, and **886**, after which normalization layers **880**, **884**, and **888**, respectively, are included.

(125) Multi-head attention layers **862**, **878**, and **882** may each linearly project queries, keys, and values using a set of linear projections to a corresponding dimension. Each linear projection may be processed using an attention function (e.g., dot-product or additive attention), thereby yielding n-dimensional output values for each linear projection. The resulting values may be concatenated and once again projected, such that the values are subsequently processed as illustrated in FIG. **8B** (e.g., by a corresponding normalization layer **864**, **880**, or **884**).

(126) Feed forward layers **866** and **886** may each be a fully connected feed-forward network, which applies to each position. In examples, feed forward layers **866** and **886** each include a plurality of linear transformations with a rectified linear unit activation in between. In examples, each linear transformation is the same across different positions, while different parameters may be used as compared to other linear transformations of the feed-forward network.

(127) Additionally, aspects of linear transformation **892** may be similar to the linear transformations discussed above with respect to multi-head attention layers **862**, **878**, and **882**, as well as feed forward layers **866** and **886**. Softmax **894** may further convert the output of linear transformation **892** to predicted next-token probabilities, as indicated by output probabilities **896**. It will be appreciated that the illustrated architecture is provided in as an example and, in other examples, any of a variety of other model architectures may be used in accordance with the disclosed aspects. In some instances, multiple iterations of processing are performed according to the above-described aspects (e.g., using generative model package **804** in FIG. **8A** or encoder **852** and decoder **854** in FIG. **8B**) to generate a series of output tokens (e.g., words), for example which are then combined to yield a complete sentence (and/or any of a variety of other content). It will be appreciated that other generative models may generate multiple output tokens in a single iteration and may thus used a reduced number of iterations or a single iteration.

(128) Accordingly, output probabilities **896** may thus form embedding output **806** according to aspects described herein, such that the output of the generative ML model (e.g., which may include structured output) is used as input for determining an action according to aspects described herein (e.g., similar to the action determination engine **514** of FIG. **5**). In other examples, embedding

output **806** is provided as generated output for updating an embedding object memory (e.g., the embedding object memory **310**).

(129) FIGS. **9-11** and the associated descriptions provide a discussion of a variety of operating environments in which aspects of the disclosure may be practiced. However, the devices and systems illustrated and discussed with respect to FIGS. **9-11** are for purposes of example and illustration and are not limiting of a vast number of computing device configurations that may be utilized for practicing aspects of the disclosure, described herein.

(130) FIG. **9** is a block diagram illustrating physical components (e.g., hardware) of a computing device **900** with which aspects of the disclosure may be practiced. The computing device components described below may be suitable for the computing devices described above, including computing device **102** in FIG. **1**. In a basic configuration, the computing device **900** may include at least one processing unit **902** and a system memory **904**. Depending on the configuration and type of computing device, the system memory **904** may comprise, but is not limited to, volatile storage (e.g., random access memory), non-volatile storage (e.g., read-only memory), flash memory, or any combination of such memories.

(131) The system memory **904** may include an operating system **905** and one or more program modules **906** suitable for running software application **920**, such as one or more components supported by the systems described herein. As examples, system memory **904** may store embedding object memory insertion engine or component **924** and/or embedding object memory retrieval engine or component **926**. The operating system **905**, for example, may be suitable for controlling the operation of the computing device **900**.

(132) Furthermore, aspects of the disclosure may be practiced in conjunction with a graphics library, other operating systems, or any other application program and is not limited to any particular application or system. This basic configuration is illustrated in FIG. **9** by those components within a dashed line **908**. The computing device **900** may have additional features or functionality. For example, the computing device **900** may also include additional data storage devices (removable and/or non-removable) such as, for example, magnetic disks, optical disks, or tape. Such additional storage is illustrated in FIG. **9** by a removable storage device **909** and a non-removable storage device **910**.

(133) As stated above, a number of program modules and data files may be stored in the system memory **904**. While executing on the processing unit **902**, the program modules **906** (e.g., application **920**) may perform processes including, but not limited to, the aspects, as described herein. Other program modules that may be used in accordance with aspects of the present disclosure may include electronic mail and contacts applications, word processing applications, spreadsheet applications, database applications, slide presentation applications, drawing or computer-aided application programs, etc.

(134) Furthermore, aspects of the disclosure may be practiced in an electrical circuit comprising discrete electronic elements, packaged or integrated electronic chips containing logic gates, a circuit utilizing a microprocessor, or on a single chip containing electronic elements or microprocessors. For example, aspects of the disclosure may be practiced via a system-on-a-chip (SOC) where each or many of the components illustrated in FIG. **9** may be integrated onto a single integrated circuit. Such an SOC device may include one or more processing units, graphics units, communications units, system virtualization units and various application functionality all of which are integrated (or “burned”) onto the chip substrate as a single integrated circuit. When operating via an SOC, the functionality, described herein, with respect to the capability of client to switch protocols may be operated via application-specific logic integrated with other components of the computing device **900** on the single integrated circuit (chip). Some aspects of the disclosure may also be practiced using other technologies capable of performing logical operations such as, for example, AND, OR, and NOT, including but not limited to mechanical, optical, fluidic, and quantum technologies. In addition, some aspects of the disclosure may be practiced within a

general purpose computer or in any other circuits or systems.

(135) The computing device **900** may also have one or more input device(s) **912** such as a keyboard, a mouse, a pen, a sound or voice input device, a touch or swipe input device, etc. The output device(s) **914** such as a display, speakers, a printer, etc. may also be included. The aforementioned devices are examples and others may be used. The computing device **900** may include one or more communication connections **916** allowing communications with other computing devices **950**. Examples of suitable communication connections **916** include, but are not limited to, radio frequency (RF) transmitter, receiver, and/or transceiver circuitry; universal serial bus (USB), parallel, and/or serial ports.

(136) The term computer readable media as used herein may include computer storage media. Computer storage media may include volatile and nonvolatile, removable and non-removable media implemented in any method or technology for storage of information, such as computer readable instructions, data structures, or program modules. The system memory **904**, the removable storage device **909**, and the non-removable storage device **910** are all computer storage media examples (e.g., memory storage). Computer storage media may include RAM, ROM, electrically erasable read-only memory (EEPROM), flash memory or other memory technology, CD-ROM, digital versatile disks (DVD) or other optical storage, magnetic cassettes, magnetic tape, magnetic disk storage or other magnetic storage devices, or any other article of manufacture which can be used to store information and which can be accessed by the computing device **900**. Any such computer storage media may be part of the computing device **900**. Computer storage media does not include a carrier wave or other propagated or modulated data signal.

(137) Communication media may be embodied by computer readable instructions, data structures, program modules, or other data in a modulated data signal, such as a carrier wave or other transport mechanism, and includes any information delivery media. The term “modulated data signal” may describe a signal that has one or more characteristics set or changed in such a manner as to encode information in the signal. By way of example, and not limitation, communication media may include wired media such as a wired network or direct-wired connection, and wireless media such as acoustic, radio frequency (RF), infrared, and other wireless media.

(138) FIG. **10** is a block diagram illustrating the architecture of one aspect of a computing device. That is, the computing device can incorporate a system (e.g., an architecture) **1002** to implement some aspects. In some examples, the system **1002** is implemented as a “smart phone” capable of running one or more applications (e.g., browser, e-mail, calendaring, contact managers, messaging clients, games, and media clients/players). In some aspects, the system **1002** is integrated as a computing device, such as an integrated personal digital assistant (PDA) and wireless phone.

(139) One or more application programs **1066** may be loaded into the memory **1062** and run on or in association with the operating system **1064**. Examples of the application programs include phone dialer programs, e-mail programs, personal information management (PIM) programs, word processing programs, spreadsheet programs, Internet browser programs, messaging programs, and so forth. The system **1002** also includes a non-volatile storage area **1068** within the memory **1062**. The non-volatile storage area **1068** may be used to store persistent information that should not be lost if the system **1002** is powered down. The application programs **1066** may use and store information in the non-volatile storage area **1068**, such as e-mail or other messages used by an e-mail application, and the like. A synchronization application (not shown) also resides on the system **1002** and is programmed to interact with a corresponding synchronization application resident on a host computer to keep the information stored in the non-volatile storage area **1068** synchronized with corresponding information stored at the host computer. As should be appreciated, other applications may be loaded into the memory **1062** and run on the mobile computing device **1000** described herein (e.g., an embedding object memory insertion engine, an embedding object memory retrieval engine, etc.).

(140) The system **1002** has a power supply **1070**, which may be implemented as one or more

batteries. The power supply **1070** might further include an external power source, such as an AC adapter or a powered docking cradle that supplements or recharges the batteries.

(141) The system **1002** may also include a radio interface layer **1072** that performs the function of transmitting and receiving radio frequency communications. The radio interface layer **1072** facilitates wireless connectivity between the system **1002** and the “outside world,” via a communications carrier or service provider. Transmissions to and from the radio interface layer **1072** are conducted under control of the operating system **1064**. In other words, communications received by the radio interface layer **1072** may be disseminated to the application programs **1066** via the operating system **1064**, and vice versa.

(142) The visual indicator **1020** may be used to provide visual notifications, and/or an audio interface **1074** may be used for producing audible notifications via the audio transducer **1025**. In the illustrated example, the visual indicator **1020** is a light emitting diode (LED) and the audio transducer **1025** is a speaker. These devices may be directly coupled to the power supply **1070** so that when activated, they remain on for a duration dictated by the notification mechanism even though the processor **1060** and/or special-purpose processor **1061** and other components might shut down for conserving battery power. The LED may be programmed to remain on indefinitely until the user takes action to indicate the powered-on status of the device. The audio interface **1074** is used to provide audible signals to and receive audible signals from the user. For example, in addition to being coupled to the audio transducer **1025**, the audio interface **1074** may also be coupled to a microphone to receive audible input, such as to facilitate a telephone conversation. In accordance with aspects of the present disclosure, the microphone may also serve as an audio sensor to facilitate control of notifications, as will be described below. The system **1002** may further include a video interface **1076** that enables an operation of an on-board camera **1030** to record still images, video stream, and the like.

(143) A computing device implementing the system **1002** may have additional features or functionality. For example, the computing device may also include additional data storage devices (removable and/or non-removable) such as, magnetic disks, optical disks, or tape. Such additional storage is illustrated in FIG. **10** by the non-volatile storage area **1068**.

(144) Data/information generated or captured by the computing device and stored via the system **1002** may be stored locally on the computing device, as described above, or the data may be stored on any number of storage media that may be accessed by the device via the radio interface layer **1072** or via a wired connection between the computing device and a separate computing device associated with the computing device, for example, a server computer in a distributed computing network, such as the Internet. As should be appreciated such data/information may be accessed via the computing device via the radio interface layer **1072** or via a distributed computing network. Similarly, such data/information may be readily transferred between computing devices for storage and use according to well-known data/information transfer and storage means, including electronic mail and collaborative data/information sharing systems.

(145) FIG. **11** illustrates one aspect of the architecture of a system for processing data received at a computing system from a remote source, such as a personal computer **1104**, tablet computing device **1106**, or mobile computing device **1108**, as described above. Content displayed at server device **1102** may be stored in different communication channels or other storage types. For example, various documents may be stored using a directory service **1124**, a web portal **1125**, a mailbox service **1126**, an instant messaging store **1128**, or a social networking site **1130**.

(146) An application **1120** (e.g., similar to the application **920**) may be employed by a client that communicates with server device **1102**. Additionally, or alternatively, embedding object memory insertion engine **1121** and/or embedding object memory retrieval engine **1122** may be employed by server device **1102**. The server device **1102** may provide data to and from a client computing device such as a personal computer **1104**, a tablet computing device **1106** and/or a mobile computing device **1108** (e.g., a smart phone) through a network **1115**. By way of example, the

computer system described above may be embodied in a personal computer **1104**, a tablet computing device **1106** and/or a mobile computing device **1108** (e.g., a smart phone). Any of these examples of the computing devices may obtain content from the store **1116**, in addition to receiving graphical data useable to be either pre-processed at a graphic-originating system, or post-processed at a receiving computing system.

(147) As will be understood from the foregoing disclosure, one aspect of the technology relates to a method for storing an entry in an embedding object memory. The method comprises: receiving a content item, the content item having one or more content data; providing one of the content data associated with the content item to one or more semantic embedding models, wherein the one or more semantic embedding models generate one or more semantic embeddings; receiving, from one or more of the semantic embedding models, one or more semantic embeddings, wherein a collection of semantic embeddings is associated with a first semantic embedding model of the one or more semantic embedding models, and wherein the collection of semantic embeddings comprises a first semantic embedding generated by the first semantic embedding model for at least one content data from the respective content item; inserting the one or more semantic embeddings into the embedding object memory, wherein the embedding object memory stores one or more semantic embeddings from the collection of semantic embeddings, wherein the one or more semantic embeddings are associated with a respective indication corresponding to a reference to source data associated with the one or more semantic embeddings, and wherein the insertion triggers a spatial storage operation to store a vector representation of the one or more semantic embeddings; and providing the embedding object memory. In some examples, the vector representation is stored in at least one of an approximate nearest neighbor (ANN) tree, a k-d tree, or a multidimensional tree. In some examples, the semantic embedding models comprise a generative large language model (LLM). In some examples, the content data are a plurality of content data that are each a respective one of audio content data, visual content data, gaze content data, weather content data, news content data, calendar content data, email content data, or location content data. In some examples, the embedding object memory is stored at a location that is different than the location of the source data. In some examples, the one or more semantic embedding models comprise a version, wherein each of the semantic embeddings generated by each of the respective models comprise metadata corresponding to the version, and wherein the method further comprises: providing an updated semantic embedding model to replace at least one of the semantic embedding models, the updated semantic embedding model comprising an updated version that is different than the version of the at least one of the semantic embedding models; receiving, from the updated semantic embedding model, an updated one or more semantic embeddings corresponding to the one or more semantic embeddings generated by the at least one of the semantic embedding models; and inserting the updated semantic embeddings in the embedding object memory with metadata corresponding to the updated version. In some examples, the method further comprises: deleting the one or more semantic embeddings corresponding to the updated semantic embeddings, the one or more semantic embeddings having metadata corresponding to a version that is different than the updated version.

(148) In another aspect, the technology relates to a method for retrieving information from an embedding object memory. The method comprises: receiving an input embedding, wherein the input embedding is generated by a machine-learning model; retrieving a plurality of collections of stored semantic embeddings, from the embedding object memory, based on the input embedding, wherein the plurality of collections of stored semantic embeddings each correspond to respective content data; retrieving a subset of semantic embeddings from at least one of the plurality of collections of stored semantic embeddings based on a similarity to the input embedding; determining, based on the subset of semantic embeddings and the input embedding, an action; and providing the action as an output. In some examples, each semantic embedding of the subset of semantic embeddings is associated with source data corresponding to the respective content data,

wherein the determining an action comprises: locating the source data; and determining the action based on the input embedding and the source data. In some examples, the source data includes one or more of audio files, text files, or image files. In some examples, the retrieving a subset of embeddings comprises: determining a respective similarity between the input embedding and each embedding of the plurality of collections of stored semantic embeddings; determining an ordered ranking of the one or more similarities or that one or more of the similarities are less than a predetermined threshold; and retrieving the subset of semantic embeddings with similarities to the input embedding that are less than the predetermined threshold or based on the ordered ranking, thereby retrieving subsets of semantic embeddings from at least one of the collections of semantic embeddings that are determined to be related to the input embedding. In some examples, the similarities are distances. In some examples, the method further comprises, prior to receiving the input embedding: receiving user-input; and generating the input embedding based on the user-input. In some examples, the method further comprises adapting a computing device to perform the action.

(149) In a further aspect, the technology relates to a method for inserting entries into and retrieving information from an embedding object memory. The memory comprises: receiving a content item, the content item having one or more content data; providing one of the content data associated with the content item to one or more semantic embedding models, wherein the one or more semantic embedding models generate one or more semantic embeddings; receiving, from one or more of the semantic embedding models, one or more semantic embeddings, wherein a collection of semantic embeddings is associated with a first semantic embedding model of the one or more semantic embedding models, and wherein the collection of semantic embeddings comprises a first semantic embedding generated by the first semantic embedding model for at least one content data from the respective content item; inserting the one or more semantic embeddings into the embedding object memory, wherein the embedding object memory stores one or more semantic embeddings from the collection of semantic embeddings, and wherein the one or more semantic embeddings are associated with a respective indication corresponding to a reference to source data associated with the one or more semantic embeddings; receiving an input embedding; retrieving a plurality of collections of stored semantic embeddings, from the semantic object memory, based on the input embedding; retrieving a subset of semantic embeddings from at least one of the plurality of collections of stored semantic embeddings based on a similarity to the input embedding; and providing the subset of semantic embeddings as an output. In some examples, the one or more semantic embedding models comprise a version, wherein each of the semantic embeddings generated by each of the respective semantic embedding models comprise metadata corresponding to the version, and wherein the plurality of collections of stored semantic embeddings are retrieved based on the metadata corresponding to the version. In some examples, the method further comprises: determining, based on the subset of semantic embeddings and the input embedding, an action; and adapting a computing device to perform the action. In some examples, the retrieving a subset of semantic embeddings comprises: determining a respective similarity between the input embedding and each semantic embedding of the plurality of collections of stored semantic embeddings; determining that one or more of the similarities are less than a predetermined threshold; and retrieving the subset of semantic embeddings with similarities to the input embedding that are less than the predetermined threshold, thereby retrieving subsets of semantic embeddings from at least one of the collections of semantic embeddings that are determined to be related to the input embedding. In some examples, the semantic embedding models comprise a generative large language model (LLM). In some examples, the content data are provided to the one or more semantic embedding models locally or via an application programming interface (API).

(150) Aspects of the present disclosure, for example, are described above with reference to block diagrams and/or operational illustrations of methods, systems, and computer program products

according to aspects of the disclosure. The functions/acts noted in the blocks may occur out of the order as shown in any flowchart. For example, two blocks shown in succession may in fact be executed substantially concurrently or the blocks may sometimes be executed in the reverse order, depending upon the functionality/acts involved.

(151) The description and illustration of one or more aspects provided in this application are not intended to limit or restrict the scope of the disclosure as claimed in any way. The aspects, examples, and details provided in this application are considered sufficient to convey possession and enable others to make and use claimed aspects of the disclosure. The claimed disclosure should not be construed as being limited to any aspect, example, or detail provided in this application. Regardless of whether shown and described in combination or separately, the various features (both structural and methodological) are intended to be selectively included or omitted to produce an embodiment with a particular set of features. Having been provided with the description and illustration of the present application, one skilled in the art may envision variations, modifications, and alternate aspects falling within the spirit of the broader aspects of the general inventive concept embodied in this application that do not depart from the broader scope of the claimed disclosure.

Claims

1. A method for storing an entry in an embedding object memory, the method comprising: receiving a content item, the content item having one or more content data; providing one of the content data associated with the content item to one or more semantic embedding models, wherein the one or more semantic embedding models generate one or more semantic embeddings; receiving, from one or more of the semantic embedding models, one or more semantic embeddings, wherein a collection of semantic embeddings is associated with a first semantic embedding model of the one or more semantic embedding models, wherein the collection of semantic embeddings comprises a first semantic embedding generated by the first semantic embedding model for at least one content data from the respective content item, wherein the one or more semantic embedding models comprise a version, and wherein each of the semantic embeddings generated by each of the respective one or more semantic embedding models comprise metadata corresponding to the version; inserting the one or more semantic embeddings into the embedding object memory, wherein the embedding object memory stores one or more semantic embeddings from the collection of semantic embeddings, wherein the one or more semantic embeddings are associated with a respective indication corresponding to a reference to source data associated with the one or more semantic embeddings, and wherein the insertion triggers a spatial storage operation to store a vector representation of the one or more semantic embeddings; providing an updated semantic embedding model to replace at least one of the semantic embedding models, the updated semantic embedding model comprising an updated version that is different than the version of the at least one of the semantic embedding models; receiving, from the updated semantic embedding model, an updated one or more semantic embeddings corresponding to the one or more semantic embeddings generated by the at least one of the semantic embedding models, wherein the updated one or more semantic embeddings are generated based on the one of the content data used to generate the one or more semantic embeddings; inserting the updated semantic embeddings in the embedding object memory with metadata corresponding to the updated version; and providing the embedding object memory.

2. The method of claim 1, wherein the vector representation is stored in at least one of an approximate nearest neighbor (ANN) tree, a k-d tree, or a multidimensional tree.

3. The method of claim 1, wherein the semantic embedding models comprise a generative large language model (LLM).

4. The method of claim 1, wherein the content data are a plurality of content data that are each a respective one of audio content data, visual content data, gaze content data, weather content data,

news content data, calendar content data, email content data, or location content data.

5. The method of claim 1, wherein the embedding object memory is stored at a location that is different than the location of the source data.

6. The method of claim 1, further comprising: deleting the one or more semantic embeddings corresponding to the updated semantic embeddings, the one or more semantic embeddings having metadata corresponding to a version that is different than the updated version.

7. A method for retrieving information from an embedding object memory, the method comprising: receiving an input embedding, wherein the input embedding is generated by a machine-learning model; retrieving a plurality of collections of stored semantic embeddings, from the embedding object memory, based on the input embedding, wherein the plurality of collections of stored semantic embeddings each correspond to respective content data; retrieving a subset of semantic embeddings from at least one of the plurality of collections of stored semantic embeddings based on a similarity to the input embedding, wherein the subset of semantic embeddings includes a first semantic embedding stored with metadata corresponding to a first version of a semantic embedding model used to generate the first semantic embedding based on a content item, and wherein the subset of semantic embeddings further includes a second semantic embedding stored with metadata corresponding to a second version of the semantic embedding model used to generate the second semantic embedding based on the content item; determining, based on the subset of semantic embeddings and the input embedding, an action; and providing the action as an output.

8. The method of claim 7, wherein each semantic embedding of the subset of semantic embeddings is associated with source data corresponding to the respective content data, wherein the determining an action comprises: locating the source data; and determining the action based on the input embedding and the source data.

9. The method of claim 8, wherein the source data includes one or more of audio files, text files, or image files.

10. The method of claim 7, wherein the retrieving a subset of embeddings comprises: determining a respective similarity between the input embedding and each embedding of the plurality of collections of stored semantic embeddings; determining an ordered ranking of the one or more similarities or that one or more of the similarities are less than a predetermined threshold; and retrieving the subset of semantic embeddings with similarities to the input embedding that are less than the predetermined threshold or based on the ordered ranking, thereby retrieving subsets of semantic embeddings from at least one of the collections of semantic embeddings that are determined to be related to the input embedding.

11. The method of claim 10, wherein the similarities are distances.

12. The method of claim 7, further comprising, prior to receiving the input embedding: receiving user-input; and generating the input embedding based on the user-input.

13. The method of claim 7, further comprising: adapting a computing device to perform the action.

14. The method of claim 7, wherein the machine-learning model is a generative multimodal machine-learning model.

15. A method for inserting entries into and retrieving information from an embedding object memory, the method comprising: receiving a content item, the content item having one or more content data; providing one of the content data associated with the content item to one or more semantic embedding models, wherein the one or more semantic embedding models generate one or more semantic embeddings; receiving, from one or more of the semantic embedding models, one or more semantic embeddings, wherein a collection of semantic embeddings is associated with a first semantic embedding model of the one or more semantic embedding models, wherein the collection of semantic embeddings comprises a first semantic embedding generated by the first semantic embedding model for at least one content data from the respective content item, wherein the one or more semantic embedding models comprise a version, and wherein each of the semantic embeddings generated by each of the respective one or more semantic embedding models comprise

metadata corresponding to the version; inserting the one or more semantic embeddings into the embedding object memory, wherein the embedding object memory stores one or more semantic embeddings from the collection of semantic embeddings, and wherein the one or more semantic embeddings are associated with a respective indication corresponding to a reference to source data associated with the one or more semantic embeddings; providing an updated semantic embedding model to replace at least one of the semantic embedding models, the updated semantic embedding model comprising an updated version that is different than the version of the at least one of the semantic embedding models; receiving, from the updated semantic embedding model, an updated one or more semantic embeddings corresponding to the one or more semantic embeddings generated by the at least one of the semantic embedding models, wherein the updated one or more semantic embeddings are generated based on the one of the content data used to generate the one or more semantic embeddings; inserting the updated semantic embeddings in the embedding object memory with metadata corresponding to the updated version; receiving an input embedding; retrieving a plurality of collections of stored semantic embeddings, from the semantic object memory, based on the input embedding; retrieving a subset of semantic embeddings from at least one of the plurality of collections of stored semantic embeddings based on a similarity to the input embedding; and providing the subset of semantic embeddings as an output.

16. The method of claim 15, wherein the plurality of collections of stored semantic embeddings are retrieved based on the metadata corresponding to the version.

17. The method of claim 15, further comprising: determining, based on the subset of semantic embeddings and the input embedding, an action; and adapting a computing device to perform the action.

18. The method of claim 15, wherein the retrieving a subset of semantic embeddings comprises: determining a respective similarity between the input embedding and each semantic embedding of the plurality of collections of stored semantic embeddings; determining that one or more of the similarities are less than a predetermined threshold; and retrieving the subset of semantic embeddings with similarities to the input embedding that are less than the predetermined threshold, thereby retrieving subsets of semantic embeddings from at least one of the collections of semantic embeddings that are determined to be related to the input embedding.

19. The method of claim 15, wherein the semantic embedding models comprise a generative large language model (LLM).

20. The method of claim 15, wherein the content data are provided to the one or more semantic embedding models locally or via an application programming interface (API).
