

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12405937
Kind Code	B1
Date of Patent	September 02, 2025
Inventor(s)	Matthews; William Brad et al.

Hash operation manipulations

Abstract

Certain hash-based operations in network devices and other devices, such as mapping and/or lookup operations, are improved by manipulating a hash key prior to executing a hash function on the hash key and/or by manipulating outputs of a hash function. A device may be configured to manipulate hash keys and/or outputs using manipulation logic based on one or more predefined manipulation values. A similar hash-based operation may be performed by multiple devices within a network of computing devices. Different devices may utilize different predefined manipulation values for their respective implementations of the manipulation logic. For instance, each device may assign itself a random mask value for key transformation logic as part of an initialization process when the device powers up and/or each time the device reboots. In an embodiment, described techniques may increase the entropy of hashing function outputs in certain contexts, thereby increasing the effectiveness of certain hashing functions.

Inventors:	Matthews; William Brad (Los Gatos, CA), Agarwal; Puneet (Santa Clara, CA)
Applicant:	Innovium, Inc. (Santa Clara, CA)
Family ID:	72664151
Assignee:	Innovium, Inc. (Santa Clara, CA)
Appl. No.:	18/754025
Filed:	June 25, 2024

Related U.S. Application Data

continuation parent-doc US 17061104 20201001 US 12019606 child-doc US 18754025
continuation parent-doc US 15359616 20161122 US 10795873 20201006 child-doc US 17061104

Publication Classification

Int. Cl.: G06F16/22 (20190101); G06F16/23 (20190101); G06F16/2455 (20190101)

U.S. Cl.:

CPC G06F16/2255 (20190101); G06F16/2379 (20190101); G06F16/2455 (20190101);

Field of Classification Search

CPC: G06F (16/2255); G06F (16/2379); G06F (16/2455)

USPC: 707/756

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
4961139	12/1989	Hong et al.	N/A	N/A
5251165	12/1992	James	N/A	N/A
6345281	12/2001	Kardos et al.	N/A	N/A
6430562	12/2001	Kardos et al.	N/A	N/A
6434662	12/2001	Greene et al.	N/A	N/A
6782475	12/2003	Sumner	N/A	N/A
7274694	12/2006	Cheng et al.	N/A	N/A
7895047	12/2010	Wetzer et al.	N/A	N/A
8285957	12/2011	Giridhar et al.	N/A	N/A
8433547	12/2012	Dalgas et al.	N/A	N/A
8874866	12/2013	Tyson	N/A	N/A
9230620	12/2015	Lee	N/A	N/A
9455966	12/2015	Roitshtein	N/A	H04L 9/0819
9716592	12/2016	Mandal et al.	N/A	N/A
10169039	12/2018	Moudgill et al.	N/A	N/A
2002/0116527	12/2001	Chen et al.	N/A	N/A
2003/0135625	12/2002	Fontes et al.	N/A	N/A
2006/0248095	12/2005	Cozzani	N/A	N/A
2008/0063187	12/2007	Yoshida et al.	N/A	N/A
2008/0181103	12/2007	Davies	N/A	N/A
2009/0085780	12/2008	Vijayrangan	N/A	N/A
2009/0100246	12/2008	Uehara	N/A	N/A
2009/0177707	12/2008	Devraj et al.	N/A	N/A
2009/0222625	12/2008	Ghosh et al.	N/A	N/A
2010/0082309	12/2009	Dawson et al.	N/A	N/A
2011/0013639	12/2010	Matthews et al.	N/A	N/A
2011/0099258	12/2010	Ayachitula et al.	N/A	N/A
2011/0235508	12/2010	Goel et al.	N/A	N/A
2012/0041569	12/2011	Zhang et al.	N/A	N/A
2012/0117231	12/2011	Ingle et al.	N/A	N/A
2012/0136999	12/2011	Roitshtein et al.	N/A	N/A
2013/0159039	12/2012	Brech et al.	N/A	N/A
2013/0173908	12/2012	Matthews et al.	N/A	N/A
2013/0347103	12/2012	Veteikis et al.	N/A	N/A
2015/0058595	12/2014	Gura et al.	N/A	N/A
2015/0120754	12/2014	Chase et al.	N/A	N/A
2015/0121035	12/2014	Steele, Jr. et al.	N/A	N/A
2015/0127900	12/2014	Dharmapurikar et al.	N/A	N/A
2015/0281082	12/2014	Rajahalme	N/A	N/A

OTHER PUBLICATIONS

U.S. Appl. No. 15/359,607, Advisory Action dated Jul. 10, 2019. cited by applicant
U.S. Appl. No. 15/359,607, Final Office Action dated Mar. 7, 2019. cited by applicant
U.S. Appl. No. 15/359,607, Non-Final Office Action dated Aug. 13, 2018. cited by applicant
U.S. Appl. No. 15/359,607, Notice of Allowance dated Nov. 14, 2019. cited by applicant
U.S. Appl. No. 15/359,609, Non-Final Office Action dated Nov. 9, 2018. cited by applicant
U.S. Appl. No. 15/359,609, Notice of Allowance dated Mar. 8, 2019. cited by applicant
U.S. Appl. No. 15/359,616, Advisory Action dated Jun. 20, 2019. cited by applicant
U.S. Appl. No. 15/359,616, Advisory Action dated Apr. 7, 2020. cited by applicant
U.S. Appl. No. 15/359,616, Final Office Action dated Apr. 18, 2019. cited by applicant
U.S. Appl. No. 15/359,616, Final Office Action dated Jan. 13, 2020. cited by applicant
U.S. Appl. No. 15/359,616, Non-Final Office Action dated Oct. 3, 2018. cited by applicant
U.S. Appl. No. 15/359,616, Non-Final Office Action dated Sep. 9, 2019. cited by applicant
U.S. Appl. No. 15/359,616, Notice of Allowance dated Jun. 4, 2020. cited by applicant
U.S. Appl. No. 16/006,073, Notice of Allowance dated Aug. 7, 2019. cited by applicant
U.S. Appl. No. 17/061,104, Advisory Action dated Dec. 14, 2023. cited by applicant

Primary Examiner: Alam; Shahid A

Attorney, Agent or Firm: Shield Intellectual Property PC

Background/Summary

PRIORITY CLAIM (1) This application claims benefit under 35 U.S.C. § 120 as a Continuation of application Ser. No. 17/061,104, filed Oct. 1, 2020, which claims benefit under 35 U.S.C. § 120 as a Continuation of application Ser. No. 15/359,616, filed Nov. 22, 2016, now U.S. Pat. No. 10,795,873, issued Oct. 6, 2020 the entire contents of which is hereby incorporated by reference as if fully set forth herein. Applicant hereby rescinds any disclaimer of claim scope in the parent application(s) or the prosecution history thereof and advise the USPTO that the claims in this application may be broader than any claim in the parent application(s).

TECHNICAL FIELD

(1) Embodiments relate generally to data lookup operations, and, more specifically, to techniques for improving hash-based lookup operations in network devices or other computer devices.

CROSS-REFERENCE TO RELATED APPLICATIONS

(2) This application is related to U.S. patent application Ser. No. 15/359,607, filed on the Nov. 22, 2016, entitled “Lens Table,” by Brad Matthews and Puneet Agarwal, the entire contents of which are hereby incorporated by reference for all purposes as if fully set forth herein. This application is related to U.S. patent application Ser. No. 15/359,609, filed on Nov. 22, 2016, entitled “Lens Distribution,” by Brad Matthews and Puneet Agarwal, the entire contents of which are hereby incorporated by reference for all purposes as if fully set forth herein.

BACKGROUND OF THE INVENTION

(3) The approaches described in this section are approaches that could be pursued, but not necessarily approaches that have been previously conceived or pursued. Therefore, unless otherwise indicated, it should not be assumed that any of the approaches described in this section qualify as prior art merely by virtue of their inclusion in this section.

(4) A hash function is any function that can be used to map data of potentially arbitrary size to data of fixed size. The mapped data, or a selected portion thereof, is said to be a hash “key.” This key is provided as input to the hash function, which then calculates an output value based on the key. The output values returned by a hash function are called hash values, or simply hashes. For a number of applications, the hash function must be deterministic—meaning that for a given input value it must always generate the same hash value. The hash function provides a mapping from one number space (or input value space) to another number space (or output value space) that is typically smaller than input value space.

(5) A hash value returned from a hash function may be utilized in a variety of computing operations, termed herein as “hash-based operations.” For example, a hash table, or hash map, is a data structure used to implement an associative array, which is a structure that maps keys to values. The hash table may be, for example, an array of buckets or slots. A hash function is utilized to compute, for a given key, an index (i.e. the hash value) into the hash table, from which a value associated with that key can be retrieved.

(6) Ideally, the hash function will assign each key to a unique bucket, but it is possible that two keys will generate an identical hash, causing both keys to point to the same bucket. Most hash table designs assume that hash collisions—different keys that are assigned by the hash function to the same bucket—will occur and must be accommodated in some way.

(7) In many situations, hash tables turn out to be more efficient than search trees or any other table lookup structure. For this reason, hash tables are widely used in many kinds of computer software, particularly for associative arrays, database indexing, caches, and sets. For instance, in networking devices, information associated with a network packet may be determined using a hash table. Various portions or fields of the packet are extracted to form a key, and this key is entered into a hash function. The resulting hash value is used to locate an index within the table at which the information associated with the packet is stored. The associated information may include, for example and without limitation, an assigned data flow, selected path, storage address, applicable policy, and/or any other suitable values.

(8) Hash functions serve yet many other purposes within computing devices, such as cryptography operations and data validation operations. Unless otherwise indicated, the techniques described herein may be utilized with respect to hash functions in any suitable computing context.

Description

BRIEF DESCRIPTION OF DRAWINGS

- (1) The present inventive subject matter is illustrated by way of example, and not by way of limitation, in the figures of the accompanying drawings and in which like reference numerals refer to similar elements and in which:
- (2) FIG. 1 is an illustrative view of various aspects of an example system in which the techniques described herein may be practiced;
- (3) FIG. 2 is an illustrative view of various aspects of an example system in which hash values may be utilized as indexes in lookup operations;
- (4) FIG. 3 is an illustrative view of various aspects of an example system in which a hash block generates multiple hash values for a single key using dual transformed keys;
- (5) FIG. 4 is an illustrative view of various aspects of an example system in which a hash block generates multiple hash values for a single key using both the original key and a transformed key;
- (6) FIG. 5 is an illustrative view of various aspects of an example system in which a hash block generates multiple hash values for a single key using a single transformed key processed by multiple hashing operations;
- (7) FIG. 6 illustrates an example transformation logic implemented within a system;
- (8) FIG. 7 illustrates a system of multiple interconnected network devices configured to utilize a same hashing function to perform a same networking function but with different input and/or output manipulations;
- (9) FIG. 8 illustrates an example flow for generating a hash value;
- (10) FIG. 9 illustrates an example flow for transforming a hash key;
- (11) FIG. 10 illustrates an example flow for using transformed keys to support example lookup operations in a network device;
- (12) FIG. 11 is an illustrative view of various aspects of an example system in which hash output manipulation techniques described herein may be practiced;
- (13) FIG. 12 is an illustrative view of various aspects of an example system in which hash values may be utilized as indexes in mapping operations;
- (14) FIG. 13 illustrates a system comprising a hash key manipulation block;
- (15) FIG. 14 illustrates an example hash output manipulation block;
- (16) FIG. 15 illustrates an example flow for transforming one or more hash values;
- (17) FIG. 16 illustrates an example flow for using transformed keys to support example mapping operations in a network device; and
- (18) FIG. 17 is block diagram of a computer system upon which embodiments of the inventive subject matter may be implemented.

DETAILED DESCRIPTION OF THE INVENTION

- (19) In the following description, for the purposes of explanation, numerous specific details are set forth in order to provide a thorough understanding of the present inventive subject matter. It will be apparent, however, that the present inventive subject matter may be practiced without these specific details. In other instances, well-known structures and devices are shown in block diagram form in order to avoid unnecessarily obscuring the present inventive subject matter.
- (20) Embodiments are described herein according to the following outline:
- (21) 1.0. General Overview
- (22) 2.0. Structural Overview 2.1. Hash Key 2.2. Transformation Logic 2.3. Input Mask Value 2.4. Hash Function 2.5. Hash Output Manipulation 2.6. Hash Key Extension 2.7. Lookup Component 2.8. Mapping Component 2.9. Example Hash Blocks 2.10. Example Transformation Logic 2.11. Example Hash Output Manipulation 2.12. Network Use Case 2.13. Miscellaneous
- (23) 3.0. Functional Overview 3.1. Generating a Hash Value 3.2. Manipulating the Hash Key 3.3. Manipulating Hash Output 3.4. Using Transformed Keys for Lookup Operations in a Network Device 3.5. Using Hash Manipulations for Mapping Functions
- (24) 4.0. Example Embodiments
- (25) 5.0. Implementation Mechanism-Hardware Overview
- (26) 6.0 Extensions and Alternatives
- (27) General Overview
- (28) Approaches, techniques, and mechanisms are disclosed for improving hash-based operations in a network device by manipulating the hash key (also referred to simply as “key”) prior to executing the hash function on the hash key and/or by manipulating one or more outputs of a hash function. For example, in some embodiments, rather than executing the hash function on the original key, the hash function is instead executed on a transformed hash key (also referred to herein simply as “transformed key”). As another example, in some embodiments, hash values computed using a hash function are manipulated before being provided to a hash-based operation that is configured to use the hash values. In some embodiments, these manipulations may be implemented by hardware-based logic, such as by a specific sequence of logic gates within a FPGA, ASIC, or other integrated circuit. In other embodiments, the manipulations may be implemented by software-based instructions executed by a general-purpose processor.
- (29) In an embodiment, a similar hash-based operation may be performed by multiple devices within a network of computing devices. For example, in a network of switches, each switch may be configured to utilize a hashing function

to look up information about or map values to packets that are being processed by the switches. To avoid certain problems that may arise in certain operating contexts if each device always hashed each key to the same value, at least some of the devices may be configured to manipulate keys and/or outputs for the hashing operation in different ways. For instance, each device may be configured to transform keys using transformation logic based on one or more predefined hash manipulation values. Different devices may thus utilize different predefined manipulation values for their respective implementations of the manipulation logic.

(30) For instance, devices may be assigned unique or random hash manipulation values for the manipulation logic when they are manufactured. Or, devices may assign themselves random hash manipulation values for the manipulation logic as part of an initialization process when the devices power up and/or each time the devices reboot. Hash manipulation values may further differ depending on specific characteristics of the hash-based operation for which the hashing is needed (e.g. a specific mapping operation, a specific lookup table upon which a lookup operation is requested, a specific group of data packets being operated upon, etc.). Hence, for example, a device may have a number of registers or other memory structures storing hash manipulation values along with information mapping those hash manipulation values to specific characteristics.

(31) In some embodiments, multiple hash values may be output for a given key. For example, a system may comprise a hash function block configured to perform multiple hashing functions. The same transformed key may be provided to each of the hashing functions. Or, some or all of the hashing functions may be associated with different transformations (e.g. using a different pre-defined mask, using a different series of logical operations, etc.), such that multiple transformed keys are generated. In yet other embodiments, one or more of the hashing functions may utilize the original key instead of a transformed key. In further embodiments, multiple hash values may be generated by submitting combinations of the original key and/or different transformed keys to the same hashing function for processing.

(32) In an embodiment, where multiple hash values are output, various configuration logic may indicate which output(s) to use in which context. In similar manner to that explained above for hash manipulation values, different characteristics of the hash-based operation for which the hash values are being generated may be mapped to different outputs, such that one output may be used to determine one type of value mapped to the key, while another output may be used to determine another type of value mapped to the key. In an embodiment, techniques described herein may be utilized in types of hashing schemes known as “cuckoo hashing.” Such schemes involve using multiple different hash functions to compute multiple possible locations in one or more hash tables at which the value associated with a key may be found.

(33) In an embodiment, an example hash key manipulation logic is as follows. The manipulation logic optionally begins with a key extension operation. For example, a number of bits corresponding to a certain predefined hash manipulation value, referred to herein as a key extension value, may be prepended or postpended to the key. The manipulation logic then proceeds to transform the key (or extended key) with a folding exclusive-or (“XOR”) operation between each sub-element (e.g. byte) of the key (or extended key) to produce a folded key. A masked addend value is then generated by performing a logical operation such as AND or XOR between the folded key and each sub-element of another predefined hash manipulation value, referred to herein as an input mask value. The masked addend is then fed to an addition function along with the original key (or extended key). The output of the addition function serves as the transformed key, which is then utilized by one or more hash functions to produce one or more hash values. In yet other embodiments, the key manipulation logic may involve other sequences of logic.

(34) In an embodiment, an example hash output manipulation logic for manipulating one or more hash values outputted by one or more hashing functions is as follows. The hash output manipulation logic begins with producing a folded key using a folding exclusive-or (“XOR”) operation between each sub-element (e.g. word) of the original key (or extended key) upon which the hash value(s) were generated. A masked addend value is then generated by performing a logical operation such as AND or XOR between the folded key and each sub-element of another predefined hash manipulation value, referred to herein as an output mask value. The output addend is then fed to a modular addition function along with the original hash value(s). The output of the modular addition function serves as the transformed hash output, which is then utilized by one or more hash-based operations to perform various mapping, lookup, or other functions. In yet other embodiments, the hash output manipulation logic may involve other sequences of logic.

(35) The described techniques and mechanisms may provide improved performance in a variety of contexts. For example, many hash-based operations involve hash functions whose input keys may be of relatively low entropy relative to each other. Such may be the case, for instance, in the context of a network device hashing packets to traffic flows. Often, packet fields such as source address, destination address, source port, protocol, and so forth are utilized to form a key. However, since many of the fields will often be limited to only a small set of possible values (e.g. because the same protocol or port predominates the network traffic), the entropy of the keys is relatively limited. Having a low entropy in keys can, with certain hash functions, limit the diversity of outputted hash values. Consequently, the set of hash values actually produced is small relative to the number of possible hash values. This means that, for example, in the context of using hash values as indexes into a lookup table, many indexes of the lookup table will never be used, thus wasting space used to store the lookup table and increasing conflicts at other indexes to which keys are frequently hashed.

(36) The techniques described herein may be utilized to improve the performance of hashing operations in these and

other contexts by manipulating the keys in manner to avoid the negative effects of low entropy. Of course, the described techniques and mechanisms may have a variety of other uses in a variety of other embodiments, and need not be limited to embodiments involving hash functions that are less effective with input keys that may be of relatively low entropy relative to each other.

(37) In other aspects, the inventive subject matter encompasses computer apparatuses and computer-readable media configured to carry out the foregoing techniques.

1.0. Structural Overview

(38) FIG. 1 is an illustrative view of various aspects of an example system **100** in which the techniques described herein may be practiced, according to an embodiment. System **100** comprises one or more computing devices. These one or more computing devices comprise any combination of hardware and/or software configured to implement the various logical components described herein, including components **110-160**. For example, system **100** may be a network device executing a sequence of logic operations implemented by one or more field-programmable gate arrays (“FPGA”), application-specific integrated circuits (“ASIC”), and/or other specialized circuitry. As another example, the one or more computing devices may include one or more memories storing instructions for implementing the various components described herein, one or more hardware processors configured to execute the instructions stored in the one or more memories, and various data repositories in the one or more memories for storing data structures utilized and manipulated by the various components.

(39) System **100** illustrates only one of many possible arrangements of components configured to provide the functionality described herein. Other arrangements may include fewer, additional, or different components, and the division of work between the components may vary depending on the arrangement.

(40) System **100** is configured to, among other functions, accept a hash key **110** as input, and generate a hash value **160** based thereon. Although only functionality for generating hash value **160** is described, it will be recognized that system **100** may in fact be configured to provide a variety of other computing functions, some of which may involve using hash values **160** generated for different keys **110** while performing various hash-based operations, and others of which may have no relation to the elements depicted in FIG. 1.

(41) 1.1. Hash Key

(42) The hash key **110** may be a value of any length (e.g. 64 bits, other numbers of bits, etc.). The key **110** is formed by or sub-dividable into a set of key elements. For example, the key elements may be a set of bytes, words, or other data units. In some embodiments, each key element is of a same length. In other embodiments, key elements may be of different lengths. In some such embodiments, the transformation logic **120** may nonetheless need to pad each element of the key **110** to a certain number of bits, so as to ensure that certain logical operations may be performed during the transformation of the key **110**.

(43) In an embodiment, the key **110** is derived from a specific unit of data for which the hash value **160** is being computed. System **100** may thus further comprise a key generator component (not depicted) configured to generate the key based on the data unit as input. For example, the data unit may comprise a set of fields, and each field in the set of fields may constitute a different key element. Key generation may comprise any suitable steps, ranging from relatively simple trimming or stripping logic that is hard-wired to generate the key by extracting values from specific bits in the data unit (e.g. the first four bytes, a value beginning at bit 5 and ending at 68, etc.), to more complex data parsing logic that interprets, decompresses, and/or decrypts the data unit to identify fields. In some embodiments, fields may furthermore be combined, divided, or processed in other ways to generate key elements.

(44) For instance, the key may be a set of fields in an Internet Protocol (“IP”) packet, Ethernet frame, or other datagram. Such fields may include, for example, a source address, destination address, source port, destination port, protocol type identifier, tag, label, or any other suitable information. As yet another non-limiting example, the key elements may be columns of a table row.

(45) A key **110** may be any suitable value used for any suitable purpose, depending on the embodiment. For example, in some embodiments, a key **110** is a value used to lookup another value, or to locate an entry in a table in which to store a value associated with the key **110**. Specific examples of such embodiments are described subsequently. In yet other embodiments, the key **110** may have other significance or functions.

(46) 1.2. Transformation Logic

(47) System **100** comprises transformation logic **120** configured to input a hash key **110** and output a transformed hash key **140**. Generally, the transformation logic **120** increases the diversity of inputs into a hashing operation **150**. The transformation logic **120** may do so using any suitable means, including shifting or transposing bits within the key **110** (including between key elements), performing logical or mathematical operations, and so forth.

(48) The transformation logic **120** may be a component taking a variety of forms depending on the embodiment. For example, in an embodiment, the transformation logic **120** may be a series of logic gates configured to perform various logic operations on the key **110**, and/or the individual elements thereof, in succession. In some embodiments, the transformation logic **120** is configured to execute quickly (e.g. a handful of clock cycles, or even one clock cycle), so as to minimize the impact of transforming the key **110** on the performance of other functions of system **100** that rely on hash value **160**. In other embodiments, the transformation logic **120** need not be as time-efficient, and may involve potentially more complex logic implemented at least in part by a general-purpose processor executing software-based

instructions.

(49) Consistency of the Transformation Logic

(50) Over the course of any continual period of normal operation of system **100**, the transformation logic must be determinative. That is, during the period, for each hashing operation in a specific context (e.g. for a specific table, set of tables, group of data units, etc.), for any given key **110**, the transformation logic **120** must always produce the same transformed key **140**, so as to preserve the consistency of any hashing operation **150** that with respect to that key **110** during that period. The transformation logic must thus be configured to always utilize the same strategy to transform a given key **110** during a period. On the other hand, in some embodiments, if normal operation is interrupted (e.g. via a reboot or restart of system **100**), transformation logic **120** may use a different strategy to produce a different transformed key **140** for a given key **110** once normal operation resumes, so long as the system **100** no longer considers any data previously computed based on the hashing operation **150** to be valid.

(51) Put another way, system **100** is considered to be in normal operation relative to the hashing function **150** any time it retains in its memory data having some aspect that is dependent on a hash value **160** calculated by hashing function **150** (e.g. a lookup table, data that has been stored in a memory structure whose address is derived from a hash value **160**, and so forth). In the event that this memory no longer stores such data (e.g. when the memory is flushed in response to the system losing power or being reset), the system is no longer considered to be in normal operation. Upon resumption of normal operation, the transformation logic **120** may thus begin utilizing a different strategy to transform keys **110**.

(52) In an embodiment, choosing a different strategy upon a system reset may be particularly useful in the event that, for example, the current strategy is not producing effective results, as may occur from time to time within certain operating contexts. A system administrator, or automatic administrative logic within the system **100**, may detect negative effects of a poorly performing strategy, and reset the unit in the hopes that selecting a new strategy will increase performance of the system **100** in the current operating context.

(53) In some embodiments, the transformed key **140** generated by transformation logic **120** is of the same length as the input key **110**. In other embodiments, the transformed key **110** may be bigger or smaller than the input key **110**.

(54) 1.3. Input Mask Value

(55) In some embodiments, transformation logic **120** may further be based upon a predefined input mask value **130**. As its name implies, the mask value **130** is used to “mask” the value of the key **110** one or more times during execution of the transformation logic **120**. For example, the transformation logic **120** may perform a logical or mathematical operation between the mask value **130** and the key **110**, individual key elements, and/or values derived therefrom. The operation may, for example, be a bitwise AND operation, addition, subtraction, multiplication, or any other suitable operation. A more specific example of transformation logic **120** using a mask value **130** is given in subsequent sections.

(56) The mask value **130** may be hard-coded into the system **100**, or set in response to various events. In the latter case, the mask value **130** may be stored in a register or any other suitable memory structure. The length of a mask value **130**, and the set of possible mask values **130**, may vary from embodiment to embodiment. In one embodiment, the mask value **130** is of the same length (e.g. number of bits) as each transformed key **140** and/or key **110**. The mask value **130** may be programmable or selected at random from a set of possible mask values. The set of possible mask values **130** may be constrained in some fashion, so as to ensure that certain mask values that have previously been identified as ineffective or undesirable are not selected.

(57) One effect of the mask value **130** is to indicate some aspect of the strategy that is used to transform each key **110**. Thus, each different mask value **130** indicates a different strategy to be utilized by the transformation logic **120**. Simply by changing the mask value **130**, a system **100** may adopt a different strategy for transforming keys **110** into transformed keys **140**.

(58) In an embodiment, to preserve the consistency of the hashing operation during normal operation of system **100**, the mask value **130** is changed only on occasions that interrupt the normal operation of system **100**, such as rebooting or resetting the system **100**, or in response to certain types of errors. This is because changing the mask value **130** will change the transformed key **140** generated for a given key **110**, and therefore invalidate any previously stored data that was derived from a hash value **160** computed prior to the change. Thus, for example, when the mask value **130** changes, lookup tables indexing keys **110** based on hash values **160** would need to be erased and re-built based on the new mask value **130**. The mask value **130** is thus considered to be “pre-defined” for a continuous period of normal operation.

(59) 1.4. Hash Function

(60) System **100** further comprises hashing logic **150** configured to accept the transformed key **140** as input, and to output a hash value **160** generated based on a hashing function of that transformed key **140**. The hashing operation may involve any suitable hashing function that accepts input values of the same character (e.g. length) as the transformed key **140** and outputs a suitable range of hash values **160** for those input values. The range of possible hash values **160** (i.e. the number of buckets to which a key **110** may hash) may be application specific, to meet the needs of the purpose for which system **100** is utilizing the hash values **160**. For example, if each hash value corresponds to a different table index or memory bucket, then the hashing function may be such as to produce a range of hash values at least approximately equal to (and no greater than) the number of available table indexes or memory buckets.

(61) The hashing operation logic **150** may be implemented by a component taking a variety of forms, depending on the

embodiment. For example, in an embodiment, the hashing operation may be implemented by a series of logic gates configured to perform various logic operations on the transformed key **140** in succession. In some embodiments, the hashing operation implements a CRC function, such as CRC-32-Ethernet or CRC-64-ISO, a Fowler-Noll-Vo (FNV) function, or other hashing function that can be implemented quickly and/or efficiently in hardware. In other embodiments, the hashing operation need not be particularly quick or efficient, and/or may be implemented by a general-purpose processor executing software-based instructions.

(62) In the case of a CRC function, any suitable polynomial may be utilized. For example, in an embodiment using a CRC-32 Ethernet function, the polynomial may be

$x^{32}+x^{26}+x^{23}+x^{22}+x^{16}+x^{12}+x^{11}+x^{10}+x^8+x^7+x^5+x^4+x^2+x+1$

(63) By implementing transformation logic **120** prior to the hash function **150**, system **100** may potentially increase the diversity of inputs and outputs of the hashing operation **150**, in certain contexts. That is, keys **110** may be highly similar and of low entropy in certain contexts (e.g. when the key elements are fields such as IP source addresses in a data center where the addresses are limited to a specific subnet). Because of this, the keys **110** may not be well-distributed throughout the range of possible inputs to the hashing operation **150**. This can in turn affect the diversity of hash values **160**, particularly in the case of certain hash functions that can be implemented highly efficiently, such as the previously-mentioned CRC operation. However, through transformation logic **120**, the keys **110** are transformed into transformed keys **140** that are much more randomly distributed throughout the range of possible input values to the hashing operation **150**.

(64) 1.5. Hash Output Manipulation

(65) FIG. **11** is an illustrative view of various aspects of an example system **1100** in which hash output manipulation techniques described herein may be practiced, according to an embodiment. System **1100** comprises one or more computing devices. These one or more computing devices comprise any combination of hardware and/or software configured to implement the various logical components described herein, including components **110-160**. For example, system **1100** may be a network device executing a sequence of logic operations implemented by one or more field-programmable gate arrays (“FPGA”), application-specific integrated circuits (“ASIC”), and/or other specialized circuitry. As another example, the one or more computing devices may include one or more memories storing instructions for implementing the various components described herein, one or more hardware processors configured to execute the instructions stored in the one or more memories, and various data repositories in the one or more memories for storing data structures utilized and manipulated by the various components.

(66) System **1100** is configured to, among other functions, accept a hash key **1110** as input, and generate transformed hash output values **1180** based thereon. Although only functionality for generating hash values **1180** is described, it will be recognized that system **1100** may in fact be configured to provide a variety of other computing functions, some of which may involve using transformed hash output values **1180** generated for different keys **1110** while performing various hash-based operations, as well as other functions of which may have no relation to the elements depicted in FIG. **11**.

(67) System **1100** comprises a hash key manipulation block **1125** configured to accept a key **1110** as input and output a transformed key **1140** based thereon. Key **1110** may be, for example, similar to key **110**. The hash key manipulation block **1125** may manipulate the key **1110** in a variety of manners. For example, the hash key manipulation block **1125** may add, prepend, or append, an extension value to the key **1110**. As another example, the hash key manipulation block **1125** may comprise transformation logic such as that described with respect to transformation logic **1120**.

(68) System **1100** further comprises a hash function block **1150**. Hash function block **1150** accepts the transformed key **1140** as input, calculates hash values based on the transformed key **1140**, and outputs those hash values as hash outputs **1160** (shown as hash outputs **1160a-n**). Each hash output **1160** may be calculated using a different hash function in hash function block **1150**. For example, hash function block **1150** may comprise a CRC-16-CCITT Ethernet function, a CRC-32C function, and a CRC-64-ISO function. In an embodiment, each hash output **1160** is output in parallel. In an embodiment, hash function block **1150** may be implemented directly by special purpose hardware configured to perform the hashing functions.

(69) The hash outputs **1160** are input into the hash output manipulation block **1170**, which generates and outputs transformed hash outputs **1180** (shown as transformed outputs **1180a-n**) based upon the inputted hash outputs **1160**. Some or all of the transformed hash outputs **1180** may then be fed as input to a hash-based operation (e.g. a lookup operation or mapping operation, not depicted).

(70) The hash output manipulation block **1170** may transform hash outputs **1160** using any suitable means, including shifting or transposing bits within the hash outputs **1160**, performing logical or mathematical operations, and so forth. The hash output manipulation block **1170** may, in some embodiments, transform the hash outputs **1160** using the same manipulation logic. In other embodiments, the hash output manipulation block **1170** may use different manipulation logic for each hash output **1160**.

(71) In an embodiment, the hash output manipulation block **1170** may transform hash outputs **1160** using a predefined output mask value **1130**. As its name implies, the mask value **1130** is used to “mask” the value of the hash output **1160** one or more times during execution of the hash output manipulation block **1170**. For example, the hash output manipulation block **1170** may perform a logical or mathematical operation between the mask value **1130** and the hash

outputs **1160**, key **1110**, and/or values derived therefrom. The operation may, for example, include a bitwise AND operation, addition, subtraction, multiplication, or any other suitable operation. A more specific example of hash output manipulation block **1170** using a mask value **1130** is given in subsequent sections. Note that predefined output mask value **1130** may have similar characteristics as predefined input mask value **130**, but need not be the same value (or even size or length).

(72) In an embodiment, the same consistency rules that apply to transformation logic **120** also apply to hash output manipulation block **1170**. Thus, for instance, if a predefined output mask value **1130** is used for a specific hashing context, the predefined output mask value **1130** for that context may only change (if at all) when there is a break in the normal operation of system **1100**, such as a system restart. Like the predefined input mask value **130**, the predefined output mask value **1130** may be assigned randomly or pseudo-randomly at such times. In yet other embodiments, the predefined output mask value is encoded in hardware and thus incapable of changing. In either case, different devices may be configured with a different predefined output mask value **1130**, and/or different predefined output mask values **1130** may be used for different hash-based operations within the same device.

(73) In an embodiment, hash outputs **1160** are concatenated together, and then subject to an operation such as modular addition with another addend value before being separated again into separate transformed outputs **1180**. The addend value may be a predefined output mask value. Or, the addend value may be computed based on a predefined mask value **1130** and some other value, such as the transformed key **1140**, key **1110**, or an extended key. Other example hash output manipulations are described elsewhere herein.

(74) The hash output manipulation block **1170** may be a component taking a variety of forms depending on the embodiment. For example, in an embodiment, the hash output manipulation block **1170** may be a series of logic gates configured to perform various logic operations on the hash outputs **1160** in succession. In some embodiments, the hash output manipulation block **1170** is configured to execute quickly (e.g. a handful of clock cycles, or even one clock cycle), so as to minimize the impact of transforming the hash outputs **1160** on the performance of other functions of system **1100** that rely on transformed hash values **1180**. In other embodiments, the manipulation logic **1170** need not be as time-efficient, and may involve potentially more complex logic implemented at least in part by a general-purpose processor executing software-based instructions.

(75) System **1100** illustrates only one of many possible arrangements of components configured to provide the functionality described herein. Other arrangements may include fewer, additional, or different components, and the division of work between the components may vary depending on the arrangement. For example, in an embodiment, hash key manipulation block **1125** may be omitted, and the hash function block **1150** is configured to accept the key **1110** directly. As another example, hash function block **1150** may produce any number of hash outputs **1160**, including just one, two, or any other number, and thus hash output manipulation block **1170** may produce the same number of transformed hash outputs **1180**.

(76) As yet another example, hash output manipulation block **1170** may be configured to deliver fewer transformed hash outputs **1180** than the hash outputs **1160** produced by hash function block **1150**, or even just one transformed hash output **1180**. For instance, hash output manipulation block **1170** may be configured to select between hash outputs **1160** based on configuration data for or additional input from a specific hash-based operation which is requesting the transformed hash outputs **1180**.

(77) 1.6. Hash Key Extension

(78) In some embodiments, hash key transformation logic, such as transformation logic **120**, may be replaced or subsumed by a hash key manipulation block, which includes other operations to manipulate the hash key instead of or in addition to those of the transformation logic. Such a hash key manipulation block may include, for example, hash key extension logic configured to append or prepend a value to the hash key.

(79) For example, FIG. **13** illustrates a system **1300** comprising a hash key manipulation block **1325**, according to an embodiment. Hash key manipulation block **1325** may, for example, be used in place of transformation logic **120**. Hash key manipulation block **1325** is further an example of hash key manipulation block **1125**.

(80) Hash key manipulation block **1325** accepts a key **1310**, such as key **110**, as input. Hash key manipulation block **1325** comprises various components configured to manipulate the inputted key **1310**, including hash key extension logic **1315**. Hash key extension logic **1315** extends the key **1310** in some manner to produce an extended key **1318** of a target size. For example, hash key extension logic **1315** may prepend or append the key **1310** with a certain number of bits, whose values are indicated by an extension value **1316**. For example, if the key **1315** is ninety bits in length, and the extended value **1318** is intended to be ninety-six bits, the extension value **1316** may be a set of six bits that are placed before or after those of the key **1310**. As another example, the extension value **1316** may be added to the key **1310**.

(81) As with transformation logic **120**, the manner by which the extension logic **1315** extends the key **1310** should be consistent over a specific context throughout normal operation of the device. Hence, if extension value **1316** is used in a specific context, the extension value **1316** should be fixed for that context throughout any continuous period of normal operation. However, if the extension value **1316** is not encoded in hardware, the extension value may be randomly assigned when the device reboots or restarts, as described with respect to the input mask value. Moreover, different devices may be configured with different extension values, and/or different extension values may be used for different

hash-based operations within the same device.

(82) The extended key **1328** is then inputted into transformation logic **1320**, such as described with respect to transformation logic **120**. An input mask value **1330**, of the same nature as mask value **130**, may be utilized by transformation logic **1320** to transform the extended key **1318** into a transformed key **1340**, which can then be used as input into any suitable hashing function.

(83) System **1300** illustrates only one of many possible arrangements of components configured to manipulate a hash key. Other arrangements may include fewer, additional, or different components, and the division of work between the components may vary depending on the arrangement.

(84) 1.7. Lookup Component

(85) FIG. **2** is an illustrative view of various aspects of an example system **200** in which hash values, such as those generated using techniques described above, may be utilized as indexes **260** in lookup operations, according to an embodiment. System **200** comprises one or more computing devices. These one or more computing devices comprise any combination of hardware and/or software configured to implement the various logical components described herein, including components **250-290**. For example, system **200** may be a network device executing a sequence of logic operations implemented by one or more field-programmable gate arrays (“FPGA”), application-specific integrated circuits (“ASIC”), and/or other specialized circuitry. As another example, the one or more computing devices may include one or more memories storing instructions for implementing the various components described herein, one or more hardware processors configured to execute the instructions stored in the one or more memories, and various data repositories in the one or more memories for storing data structures utilized and manipulated by the various components.

(86) In an embodiment, the components of systems **100** and/or **1100** may be utilized together with system **200**, and thus systems **100** and/or **1100** may be part of a same system as system **200**. More specifically, in implementing a lookup operation, system **200** illustrates one of the many possible applications of the hash values **160** generated by the components of system **100** or transformed hash values **1180** generated by the components of system **1100**. However, systems **100** and **1100** need not necessarily be utilized only in conjunction with system **200**. Note that system **200** illustrates only one of many possible arrangements of components configured to provide the lookup functionality described herein. Other arrangements may include fewer, additional, or different components, and the division of work between the components may vary depending on the arrangement.

(87) System **200** comprises a lookup component **290** configured to, among other aspects, search for keys **210** at indexes **260** within one or more tables **270**. Though the term “table” is used for convenience, the term table as used herein should more generally be understood to be any data structure representing an associative array that associates buckets with indexes. A table **270** comprises a plurality of buckets **276**, each associated with a different index number **274**. Certain buckets **276** are empty, in that they do not (yet) store any data. Each of the non-empty buckets **276** stores an entry specifying an association between a different key **210** and its corresponding lookup value. The lookup value may be any value that is stored in association with the key, such as a memory address, a path label, a priority indicator, or any other suitable value. In some embodiments, no lookup value is stored, but rather the mere existence of the key in a bucket **276** signifies an implicit lookup value.

(88) To perform a lookup operation, lookup component **290** is configured to receive a search key **210** as input. The search key **210** may be, for example, a key **110** or **1110**. The search key **210** is also sent to a hash block **250**. The hash block **250** is configured to identify one or more indexes (e.g. the depicted indexes **260a** and **260b**) for the search key **210**. The indexes **260** represent locations within the one or more tables **270** in which the key **210** may be stored. The hash block **250** computes these indexes **260** using a variety of techniques. For example, at least one of the indexes may be a hash value **160** or a transformed hash value **1180**, or a remainder when such a hash value **160** or a transformed hash value **1180** is divided by the total number n of buckets **276** that are found in the table **270**.

(89) Each index **260** should be computed using a different technique (e.g. different transformed hash key, different hash function, different output manipulation, etc.). Though two indexes **260** are shown, the number of outputs of the hash block **250** that are computed and/or utilized depends on the embodiment. Example hash blocks **250** are described subsequently.

(90) The lookup component **290** locates the index **274** that corresponds to a given input index **260**, and determines whether the bucket **276** stored at that index **274** includes the search key **210**. If so, the lookup component outputs a lookup result **280** that includes the lookup value that the bucket associates with that search key **210**.

(91) Where more than one index **260** is inputted into the lookup operation **290**, each index **260** represents an alternate index. If the search key **210** is not found at one index **260**, the lookup component **290** is configured to search the other alternate index(es) **260**, in turn, until the search key is found **210**, or until all indexes **260** have been searched. This strategy of searching more than one index is an aspect of a scheme referred to as “cuckoo” hashing. In the depicted embodiment, there are separate tables **270a** and **270b** for each of the generated indexes **260a** and **260b**, referred to as a left table and right table, respectively. The tables **270a** and **270b** have their own separate indexes **274a** and **274b** with corresponding buckets **276a** and **276b**, respectively. Index **260a** is used to locate the corresponding index **274a** in table **270a**, while index **260b** is used to locate the corresponding index **274b** in table **270b**. Although two tables **270** are depicted, there may be any number of tables **270**. For example, in some embodiments, there may be three, four, or even

more tables **270**. In other embodiments, there may be only one table **270**, and indexes **260a** and **260b** are both used to locate indexes **274** within that one table **270**.

(92) In some embodiments, only one association is stored in a bucket **276**. In other embodiments, to increase flexibility in cases where multiple keys **210** hash to the same index **260**, a bucket **276** may be a data structure capable of storing associations for more than one key **210**. Such a bucket **276** is said to store multiple entries, each corresponding to a different association. The lookup component **290** compares the search key **210** to each association stored in the bucket **276**, and, if found the search key **210** is found in the associations, returns the lookup value for the association whose key matches the search key **210**. The maximum number of entries stored per bucket **276** may be limited to some number (e.g. four, two, etc.), which may be chosen based on factors such as the character of the hash block **270** used to calculate indexes **260**, memory constraints, and so forth. The number of entries to store per bucket **276** is generally chosen in such a way as to balance the cost of memory allocated to store those entries versus the performance gained from having more available entries.

(93) Although system **200** specifically illustrates a lookup operation involving two alternate indexes **260**, it will be apparent that the techniques scale to any number of alternate indexes **260**, and may furthermore involve any number of tables **270**. Moreover, in embodiments where cuckoo hashing is not implemented, system **200** may instead be implemented using a single index **260** produced by a hash block comprising a single hash function, and using a single table **270**, without searching for alternate indexes.

(94) In an embodiment, a system **200** may comprise many lookup components **290**, each operating on their own respective tables. In another embodiment, lookup component **290** may be general-purpose, in that it searches a different set of one or more tables depending on a table set or type of value for which the lookup operation is being requested. Hence, a lookup component **290** would search a first set of tables **270** for a lookup operation to locate an assigned data flow, and a second (and different) set of tables **270** for a lookup operation to locate an assigned queue. The lookup component **290** may further be configured to provide input into the hash block **250** that indicates, for instance, a number of buckets to hash to, or other characteristics of the hashing operation or transformation logic to be performed by the hash block **250**, since different types of tables **270** may have different characteristics.

(95) In an embodiment, a pre-defined input mask value to use in the hash block may be stored in configuration data associated with the set of tables that is being searched by the key. Hence, the hash block may be configured to read or receive the mask value from different configuration data, depending on the target of the lookup operation for which a hash function is being utilized.

(96) Inserting Entries to the Lookup Table

(97) In some embodiments, lookup component **290** or another component within system **200** may further be configured to write a key **210** and its associated data to tables **270**. Subsequently, when the same key **210** is received (e.g. because a new data unit that is being processed has the same key **210**), the entry storing the key **210** may be located and the associated data returned (e.g. for processing the new data unit). For instance, when a new key **210** is received, a component of system **200** may be configured to insert its corresponding value to one of tables **270** at an index **260** calculated by hash block **250**, which index **260** is effectively the hash value or hash output of the hash block **250**. Or, if no bucket **276** at any of the computed indexes **260** contains a search key **210** for which a lookup operation is requested, the search key **210** may be added to a bucket **276** at one of the corresponding indexes **274**.

(98) A specific index **274** at which to insert the search key **210** may be determined using a variety of techniques. For instance, if the bucket **276** associated with index **260a** is already filled with the maximum number of associations, and the bucket **276** associated with index **260b** is not filled, index **260b** may be chosen (and vice versa). The search key **210** may be inserted at a specific one of the indexes **260** by default, unless the corresponding bucket **276** is already filled. For instance, index **260a** may be considered a primary index **260** to which the search key **210** should be written by default, while index **260b** is considered a secondary index **260**. One write policy may be to attempt to write a key **210** that is currently not found at the primary or secondary index to the primary index **260** first, and to the secondary index **260** if the primary index **260** is full. However, another policy may be to attempt to write the key **210** to the secondary index first. Yet another policy would be to attempt to insert the key **210** to a randomly selected one of the primary index **260** and secondary index **260** first, and to the other second. Such policies are easily extended to three, four, or even more tables **270**.

(99) If a lookup value is to be stored with the key **210**, the component updating the lookup table may be configured to receive or even request this value. For example, in the case of a write operation, the write operation may specify both the key **210** and its lookup value. As another example, where the lookup value is a memory address, an available address may be requested from a memory allocation system prior to adding the key **210** to a lookup table **270**. The memory address may then be returned as part of the operation.

(100) If none of the buckets **276** associated with the calculated indexes **260** have available space in which to store the key **210**, the component writing the search key **270** may be configured to move or delete an association already stored within those buckets **276**, in certain embodiments. Because the cuckoo hashing scheme yields at least two different locations in which a given key may be stored, as long as the set of tables **270** associated with the lookup operation are suitably sized, room may generally be made available to store most keys **210** whose assigned locations are full by moving existing keys (potentially recursively) to other locations in which the cuckoo hashing scheme permits them to

be stored. Both the keys and their associated data are moved. On the other hand, in some embodiments, such movement may introduce undesirable overhead, and the system **200** may simply be configured in at least some circumstances to delete an existing key to make room for a new key when needed, as opposed to moving the existing key. Or, system **200** may simply drop the new key without writing the new key to a table **270**.

(101) For instance, if a first key hashes to an index A and an index B, but index A contains a second key and index B contains a third key (and system is limited to one key per bucket), then the system may choose to move or delete the second key or third key. The specific key chosen for movement or deletion may depend on a variety of factors, such as which key is stored at the primary index for the first key, which key is oldest, or any other suitable factor. If the chosen key is to be moved rather than deleted, then the chosen key is resubmitted to the lookup operation. The chosen key may hash, in turn, to an index A (or B) and an index C. The chosen key may thus be moved to index C. If index C is also unavailable, then, depending on the embodiment, the key that was not initially chosen for movement may be moved instead, the chosen key may simply be deleted, or the key that is stored at index C may be moved or deleted to make room for the chosen key. In the latter case, this movement may be repeated over multiple iterations until an index has been found to store each key. Or, certain constraints may be imposed on the number of movements before a key insertion attempt is determined to have failed. The failed key may be the first key, or a key that was chosen for movement, depending on the embodiment. In an embodiment, the key that is being inserted is generally the key that is determined to have failed, and is thus not inserted.

(102) The foregoing techniques extend to embodiments where multiple keys may be stored per bucket, in which case keys need only be moved if the maximum number of keys have been stored in each bucket to which a key can be added or moved.

(103) 1.8. Mapping Component

(104) FIG. **12** is an illustrative view of various aspects of an example system **1200** in which hash values, such as those generated using techniques described above, may be utilized as indexes **1260** in mapping operations, according to an embodiment. System **1200** comprises one or more computing devices. These one or more computing devices comprise any combination of hardware and/or software configured to implement the various logical components described herein, including components **1250-1290**. For example, system **1200** may be a network device executing a sequence of logic operations implemented by one or more field-programmable gate arrays (“FPGA”), application-specific integrated circuits (“ASIC”), and/or other specialized circuitry. As another example, the one or more computing devices may include one or more memories storing instructions for implementing the various components described herein, one or more hardware processors configured to execute the instructions stored in the one or more memories, and various data repositories in the one or more memories for storing data structures utilized and manipulated by the various components.

(105) In an embodiment, the components of systems **100** and/or **1100** may be utilized together with system **1200**, and thus systems **100** and/or **1100** may be part of a same system as system **1200**. More specifically, in implementing a mapping operation, system **1200** illustrates one of the many possible applications of the hash values **160** generated by the components of system **100** or transformed hash values **1180** generated by the components of system **1100**. However, systems **100** and **1100** need not necessarily be utilized only in conjunction with system **1200**. Note that system **1200** illustrates only one of many possible arrangements of components configured to provide the lookup functionality described herein. Other arrangements may include fewer, additional, or different components, and the division of work between the components may vary depending on the arrangement.

(106) System **1200** comprises a mapping component **1290** configured to, among other aspects, find values **1276** that are mapped to keys **1210** at indexes **1260** within a table **1270**. Table **1270** comprises a plurality of values **1276**, each associated with a different index number **1274**.

(107) To perform a mapping operation, mapping component **290** is configured to receive a search key **1210** as input. The search key **1210** may be, for example, a key **110** or **1110**. The search key **1210** is also sent to a hash block **1250**.

(108) The hash block **250** is configured to identify at least one index **1260** for the search key **1210**. The index **1260** represents a location within the table **1270** at which the value mapped to the key **1210** is stored. The hash block **1250** computes the index **1260** using a variety of techniques. For example, the index **1260** may be a hash value **160** or a transformed hash value **1180**, or a remainder when such a hash value **160** or a transformed hash value **1180** is divided by the total number n of entries found in the table **1270**.

(109) In embodiments where the hash block **1250** produces more than one output, such as in system **1100** where three different transformed hash outputs **1180** are generated, the mapping component **1290** may be configured to select which output to input as the index **1260** based on configuration data associated with the mapping component **1290** and/or the table **1270**.

(110) The mapping component **1290** locates the index **1274** that corresponds to a given input index **1260**, and returns the value **1276** stored at that index **1274** as the mapped value **1280**. Although an embodiment that relies upon a table data structure is depicted, it will be understood that in other embodiments, the data structure may take other forms, such as a list associating sets or ranges of indexes **1274** with specific values **1276**.

(111) Weighted Cost Multipathing Example

(112) The mapped value **1276** may have a variety of significances, depending on the embodiment. For example, in an

embodiment, the table **1270** may be a Weighted Cost Multipath (WCMP) table or similar structure for determining a path over which to send a network packet to a specific destination. In general, WCMP involves assigning a weight to each possible path for a destination (the destination being either a single destination node or a group of nodes such as a subnet). The technique used to map paths to network packets (via their respective keys **1210**) utilizes these weights to ensure that the probability of a packet being assigned to a given path is approximately proportional to the weight of that path relative to the weights of the other paths to the destination. For instance, a path with a weight of two might be selected twice for every time a path with a weight of one is selected. Each possible path is assigned a number of entries within the table **1270**, in proportion with its weight, and each value **1276** thus represents a path. The network packet is then sent along a path corresponding to the mapped value **1280** returned by the mapping component **1290**.

(113) Since a network device may be capable of sending packets to many destinations, the device may have many different WCMP tables **1270**. The mapping component **1290** may further be configured to determine in which table **1270** to look for the mapping based upon input data indicating the destination of the packet whose key **1210** is being submitted. In an embodiment, this table **1270** may store its own configuration data that indicates various aspects of the operations to be performed by hash block **1250**. For example, the configuration data for a table **1270** may indicate which mask values and/or extension value to utilize for the hash block **1250** when looking for the mapping stored in the table **1270**. As another example, the configuration data may indicate which transformed hash output **1180** to accept as the index **1260**. In another embodiment, such configuration data is specified on a per-group or per-flow basis, independently of a WCMP table **1270**.

(114) In an embodiment, a portion of the hash output is used to perform a lookup into a group table (e.g. WCMP group table) which provides the (a) number of members in the group and (b) base address into a member table. Then, a second lookup is performed into a member table using the index=group_base_address+(hash output mod number of group members).

(115) Of course, WCMP is but one example mapping operation. Other operations may map data packets to other types of values, or map yet other types of data units to other types of values. Other non-limiting examples include link aggregation, trunking, or other aggregation operations in a network device.

(116) 1.9. Example Hash Blocks

(117) As used herein, the term hash block, such as hash block **250**, is any collection of components that accepts a key as input and generates one or more hash values (or transformed hash values) as outputs. According to an embodiment, a simple example of a hash block is the combination of components **120-150** of system **100**. Another example of a hash block is the combination of components **1125** through **1170** of FIG. **11**. A hash block may be utilized to support any suitable hash-based operation. More complicated examples of hash blocks, featuring more than one output, are also described herein. Such hash blocks may also support any suitable hash-based operations, such as lookup operations based on cuckoo hashing. In such hash-based operations, some or all of the outputs of the hash block may be utilized for the hash-based operation.

(118) In an embodiment, the same hash block hardware may be re-used for any number of different hash-based operations within the same device, such as for different lookup operations (e.g. a lookup operation for a data flow versus a lookup operation for a cache entry). Where the hash block that generates multiple outputs is utilized for multiple hash-based operations, configuration data associated with the hash-based operation may indicate which of the outputs to utilize, the number of buckets to hash to, and so forth. For example, where the hash block is utilized for multiple different lookup operations, the lookup tables specific to a given lookup operation may include a configuration block that indicates which of the different outputs to utilize for the lookup operation.

(119) FIGS. **3-5** illustrate various aspects of example systems **300**, **400**, and **500** in which multiple hash values are generated for a single key, according to embodiments. System **300**, **400**, and **500** comprises one or more computing devices. These one or more computing devices comprise any combination of hardware and/or software configured to implement the various logical components described herein. For example, systems **300**, **400**, and **500** may be implemented by one or more field-programmable gate arrays ("FPGA"), application-specific integrated circuits ("ASIC"), and/or other specialized circuitry. As another example, systems **300**, **400**, and **500** may be implemented using components that include one or more memories storing instructions for implementing the various components described herein, one or more hardware processors configured to execute the instructions stored in the one or more memories, and various data repositories in the one or more memories for storing data structures utilized and manipulated by the various components.

(120) In various embodiments, the components of systems **300**, **400**, or **500** may be utilized together with the components of other systems described herein, and thus systems **300**, **400**, or **500** may be part of a same system as those systems. However, in other embodiments, each of systems **300**, **400**, and **500** may be utilized separately from other systems described herein.

(121) Multiple Transformed Keys

(122) FIG. **3** is an illustrative view of various aspects of an example system **300** in which a hash block **305** generates multiple hash values **360** for a single key **310** using multiple transformed keys **340**, according to an embodiment. In some embodiments, the hash block **305** of system **300** may be utilized as the hash block **250** of system **200**, with key **310** corresponding to key **210**, and hash values **360a** and **360b** corresponding to indexes **260a** and **260b**, respectively. In

other embodiments, hash block 305 may generate multiple hash values 360 for purposes other than lookup operations, and need not be limited to use in conjunction with system 200 or even system 100.

(123) A hash key 310, such as key 110 or search key 210, is provided to the hash block 305. More specifically, the key 310 is inputted into both transformation logic 320a and transformation logic 320b. Each of transformation logic 320a and 320b are components configured to transform the key 310 using techniques such as described elsewhere herein. For instance, each of transformation logic 320a and 320b may be instances of transformation logic 120. However, each of transformation logic 320a and 320b is configured to transform the key 310 in a different manner, thereby producing different transformed keys 340a and 340b, respectively. For instance, in embodiments where both of transformation logic 320a and 320b rely upon a mask value 330, such as mask value 130, the mask value 330a input into transformation logic 320a may be different than the mask value 330b input into transformation logic 320b.

(124) The transformed keys 340a and 340b are then input into hash functions 350a and 350b, respectively. Hash functions 350a and 350b are configured to output hash values 360a and 360b, respectively, which are generated based on applying hashing functions to the respectively inputted transformed keys 340a and 340b. The hash functions 350a and 350b may be, for example, separate instances of hash function 150. The hash functions 350a and 350b may both utilize different hash functions, or they may utilize a same hash function, in which case the hash functions 350a and 350b may in some embodiments actually be implemented by a same component that is configured to process transformed keys 340a and 340b separately. The hash values 360 generated are generally of the same character as hash value 160, described previously.

(125) Although only two transformed keys 340 are depicted, in yet other embodiments, there may be three or even more transformed keys 340 used to calculate different hash values 360.

(126) Concurrent Use of Key and Transformed Key

(127) FIG. 4 is an illustrative view of various aspects of an example system 400 in which a hash block 405 generates multiple hash values 460 for a single key 410 using both the original key 410 and a transformed key 440, according to an embodiment. In some embodiments, the hash block 405 of system 400 may be utilized as the hash block 250 of system 200, with key 410 corresponding to key 210, and hash values 460a and 460b corresponding to indexes 260a and 260b, respectively. In other embodiments, hash block 405 may generate multiple hash values 460 for purposes other than lookup operations, and need not be limited to use in conjunction with system 200 or even system 100.

(128) A hash key 410, such as key 110 or search key 210, is provided to the hash block 405. More specifically, the key 410 is inputted into both transformation logic 420 and a hash function 450b. Transformation logic 420 is a component configured to transform the key 410 into a transformed key 440 using techniques such as described elsewhere herein. For instance, transformation logic 420 may be an instance of transformation logic 120. Optionally, transformation logic 420 may be based on a predefined mask value 430, such as described with respect to mask value 130. The transformed key 440 is input into hash function 450a.

(129) Hash functions 450a and 450b are configured to output hash values 460a and 460b, respectively, which are generated based on applying hashing functions to the inputted transformed key 440 and key 410, respectively. The hash functions 450a and 450b may be, for example, separate instances of hash function 150. The hash functions 450a and 450b may both utilize different hash functions, or they may utilize a same hash function, in which case the hash functions 450a and 450b may in some embodiments actually be implemented by a same component that is configured to process transformed key 440 and key 410 separately. The hash values 460 generated are generally of the same character as hash value 160, described previously.

(130) According to an embodiment, a hash block may comprise the components of hash block 405, wherein the transformed key 440 and key 410 are both submitted to each of hash functions 450a and 450b, and wherein hash functions 450a and 450b implement different hashing function. Hence, a total of four different hash values are generated. For instance, suppose hash function 450a implements a CRC-32C hash function and hash function 450b implements a CRC-32 Ethernet hash function. In such an embodiment, outputs would be generated for the original key with the CRC-32C hash function, the original key with the CRC-32 Ethernet hash function, the transformed key with the CRC-32C hash function, and the transformed original key with the CRC-32 Ethernet hash function. In embodiments where the hash values may be utilized as indexes in lookup tables, each output may be utilized as an index. Or, where the hash block is utilized for multiple different lookup operations, the lookup tables specific to a given lookup operation may include a configuration block that indicates which of the different outputs to utilize for the lookup operation. Another variation on the foregoing may involve generating different transformed keys 440 using different mask values 430 for each hash function.

(131) Single Transformed Key with Multiple Hashing Operations

(132) FIG. 5 is an illustrative view of various aspects of an example system 500 in which a hash block 505 generates multiple hash values 560 for a single key 510 using a single transformed key processed by multiple hashing operations 550, according to an embodiment. In some embodiments, the hash block 505 of system 500 may be utilized as the hash block 250 of system 200, with key 510 corresponding to key 210, and hash values 560a and 560b corresponding to indexes 260a and 260b, respectively. In other embodiments, hash block 505 may generate multiple hash values 560 for purposes other than lookup operations, and need not be limited to use in conjunction with system 200 or even system 100.

(133) A hash key **510**, such as key **110** or search key **210**, is provided to the hash block **505**. More specifically, the key **510** is inputted into transformation logic **520**. Transformation logic **520** is a component configured to transform the key **510** into a transformed key **540** using techniques such as described elsewhere herein. For instance, transformation logic **520** may be an instance of transformation logic **120**. Optionally, transformation logic **520** may be based on a predefined mask value **530**, such as described with respect to mask value **130**.

(134) The transformed key **540** is then input into hash functions **550a** and **550b**, respectively. Hash functions **550a** and **550b** are configured to output hash values **560a** and **560b**, respectively, which are generated based on applying hashing functions to the transformed keys **540**. The hash functions **550a** and **550b** may be, for example, separate instances of hash function **150**. The hash functions **550a** and **550b** utilize different hash functions. The hash values **360** generated are generally of the same character as hash value **160**, described previously.

(135) Note that systems **300**, **400**, and **500** illustrate only some of many possible arrangements of components configured to implement a hash block. Other arrangements may include fewer, additional, or different components, and the division of work between the components may vary depending on the arrangement. Moreover, other embodiments may involve more than two hash functions and/or transformed keys. As yet another non-limiting example, a hash block may generate three (or more) different outputs from a same hash function using the original key along with multiple different transformed keys as input.

(136) In some embodiments, each of the example hash blocks described above may be modified to include a hash key extension block immediately before each transformation logic, or as the key is inputted into the hash block. In some embodiments, each of the example hash blocks described above may be modified to include a hash output manipulation block, such as hash output manipulation block **1470**, prior to the final output of the hash block.

(137) 1.10. Example Transformation Logic

(138) FIG. **6** illustrates an example transformation logic **620** implemented within a system **600**, according to an embodiment. The example transformation logic **620** is one example of the transformation logic described elsewhere in this disclosure, including transformation logic **120**, **320**, **420**, and **520**. The transformation logic **620** may involve, for instance, executing a sequence of logic operations implemented by one or more field-programmable gate arrays ("FPGA"), application-specific integrated circuits ("ASIC"), and/or other specialized circuitry. As another example, the transformation logic may be implemented using one or more memories storing instructions for implementing the various components described herein, one or more hardware processors configured to execute the instructions stored in the one or more memories, and various data repositories in the one or more memories for storing data structures utilized and manipulated by the various components.

(139) A hash key **610** is input into transformation logic **620**. The hash key may be an original key, such as key **110**, or an extended key, such as key **1318**. The hash key **610** comprises a set of key elements **611**, labeled k.sub.0 through k.sub.n. In an embodiment, each key element **611** is a single byte. In other embodiments, the elements **611** of key **610** may take any other suitable form, as discussed with respect to key **110**.

(140) The transformation logic folds these key elements **611** together into a folded key using an XOR-folding operation **622**. The XOR-folding operation **622** involves performing XOR operations between each key element **611**. A suitable XOR-folding operation **622** is described as follows: k.sub.0 XOR k.sub.1 XOR k.sub.2 XOR . . . k.sub.n. Note that the "folded" value produced by operation **622** is a single value of the same length as the longest individual key element. Thus, for example, if each key element **611** is eight bits, then the result of operation **622** is a single eight-bit value.

(141) For example, a first XOR operation may be performed between element k.sub.0 and element k.sub.1. A second XOR operation may be performed between the result of this first XOR operation and element k.sub.1. A third XOR operation may be performed between the result of this second XOR operation and element k.sub.2. This process may be repeated until an XOR operation has been performed with each key element **611**, with each subsequent XOR operation performed on the result of the previous XOR operation. Of course, those skilled in the art will recognize that there are a variety of arrangements of logic gates capable of producing a suitably similar folded result.

(142) The folded key is then inputted into a multi-mask function **624**. The multi-mask function **624** further inputs a fixed mask value **630**, such as the pre-defined mask value **130**. The fixed mask value **630** is sub-dividable into a set of mask elements **631**, labeled m.sub.0 through m.sub.n. The number of mask elements **631** is the same as the number of key elements **611**, and each mask element **631** is of the same length as a corresponding key element **611**. For instance, in an embodiment where each key element **611** is a byte, each mask element **631** would also be a byte.

(143) The multi-mask function **624** involves producing a temporary addend value **626** comprising a set of addend elements **627**, labeled a.sub.0 through a.sub.n. The number of addend elements **627** is the same as the number of key elements **611** and mask elements **631**, and each addend element **627** is of the same length as a corresponding mask element **631**. For instance, in an embodiment where each key element **611** is a byte, each addend element **627** would also be a byte.

(144) The multi-mask function **624** produces the addend value **626** by, for each mask element **631**, performing a logical operation between the mask element **631** and the folded key to produce a corresponding one of the addend elements **627**. The logical operation may be, for instance, an AND operation or XOR operation. For example, if the chosen logical operation is AND, the multi-mask function **624** is described as follows: a.sub.i=XOR.sub.result & m.sub.i, where i is the order of the element in the corresponding set and XOR.sub.result is the folded key. Thus, for instance, an

AND operation between m.sub.0 and the folded key would produce a.sub.0.

(145) The addend value **626** is then input into addition function **628**, along with the original key **610**. The addition function **628** performs an addition operation between the key **610** and the addend value **626** to produce a transformed key **640**. In an embodiment, transformed key **640** may be sub-dividable into a set of elements **641**, labeled x.sub.0 through x.sub.n. The number of transformed key elements **641** is the same as the number of key elements **611** and mask elements **631**, and each transformed key element **627** is of the same length as a corresponding mask element **631**

(146) In an embodiment, instead of simply adding addend **626** and input hash key **610**, addition may be performed between the individual elements of the addend **626** and key **610**. That is, $x_{sub.i} = a_{sub.i} + k_{sub.i}$. In such an embodiment, the addition of any addend element and key element may be such that the operation does not impact any other addition operations. For example, if $k_{sub.1} + a_{sub.1}$ carries over a 1, the 1 would simply be dropped rather than added to the next element (e.g. $k_{sub.2} + a_{sub.2}$). Other rules may also be imposed, such as a requirement that the result not saturate, and a requirement that the addition operation must simply roll over. In an embodiment where each element **641** is a byte, the addition operation is a bitwise addition operation.

(147) For example, consider a key of two bytes, k.sub.0 and k.sub.1, set to 0xFE and 0xFD, respectively. Suppose that a.sub.0 and a.sub.1 are 0x03 and 0x06, respectively. After a bitwise addition operation subject to all of the above constraints, $x_{sub.1} = a_{sub.1} + k_{sub.1} = 0xFE + 0x03 = 0x02$, while $x_{sub.1} = a_{sub.1} + k_{sub.1} = 0xFD + 0x06 = 0x03$. Hence, the key 0xFEFD is transformed to 0x0203.

(148) System **600** illustrates only one of many possible arrangements of components configured to transform a hash key **610** into a transformed key **640**. Other arrangements may include fewer, additional, or different components, and the division of work between the components may vary depending on the arrangement.

(149) 1.11. Example Hash Output Manipulation

(150) FIG. **14** illustrates an example hash output manipulation block **1470** implemented within a system **1400**, according to an embodiment. The example hash output manipulation block **1470** is one example of the hash output manipulation logic described elsewhere in this disclosure, including hash output manipulation block **1170**. The hash output manipulation block **1470** may involve, for instance, executing a sequence of logic operations implemented by one or more field-programmable gate arrays ("FPGA"), application-specific integrated circuits ("ASIC"), and/or other specialized circuitry. As another example, the hash output manipulation block may be implemented using one or more memories storing instructions for implementing the various components described herein, one or more hardware processors configured to execute the instructions stored in the one or more memories, and various data repositories in the one or more memories for storing data structures utilized and manipulated by the various components.

(151) A hash key **1410** is inputted into the hash output manipulation block **1470**. Hash key **1410** may be an original key, such as key **110**, or an extended key, such as extended key **1318**. The hash key **1410** comprises, or is sub dividable into, portions referred to herein as a set of key elements **1411**, labeled e.sub.0 through e.sub.n. In an embodiment, each key element **1411** is a word (i.e. a unit of sixteen bits). In other embodiments, the elements **1411** of key **1410** may take any other suitable form. Note that, where system **1400** also comprises transformation logic such as transformation logic **620**, elements **611** and elements **1411** may or may not be the same. For instance, a key **1410** that is subdivided into byte-size elements for transformation upon input may be subdivided into larger (e.g. word) sized elements **1411** for hash output manipulation block **1470**.

(152) The transformation logic folds these key elements **1411** together into a folded key using an XOR-folding operation **1422**. The XOR-folding operation **1422** involves performing XOR operations between each key element **1411**. A suitable XOR-folding operation **1422** is described as follows: $e_{sub.0} \text{ XOR } e_{sub.1} \text{ XOR } e_{sub.2} \text{ XOR } \dots \text{ XOR } e_{sub.n}$. Note that the "folded" value produced by operation **1422** is a single value of the same length as the longest individual key element **1411**. Thus, for example, if each key element **1411** is sixteen bits, then the result of operation **1422** is a single sixteen-bit value.

(153) For example, a first XOR operation may be performed between element co and element e.sub.1. A second XOR operation may be performed between the result of this first XOR operation and element e.sub.1. A third XOR operation may be performed between the result of this second XOR operation and element e.sub.2. This process may be repeated until an XOR operation has been performed with each key element **1411**, with each subsequent XOR operation performed on the result of the previous XOR operation. Of course, those skilled in the art will recognize that there are a variety of arrangements of logic gates capable of producing a suitably similar folded result.

(154) The folded key is then inputted into an addend constructor **1424**. The addend constructor **1424** further inputs a fixed output mask value **1430**, such as mask value **1130**. The fixed mask value **1430** is sub-dividable into a set of mask elements **1431**, labeled m.sub.0 through m.sub.x. Each mask element **1431** is of the same length as a corresponding key element **1411**. For instance, in an embodiment where each key element **1411** is a word, each mask element **1431** would also be a word. However, the number x of mask elements **1431** need not be the same as the number n of key elements **1411**. In fact, in an embodiment, the number x of mask elements **1431** is a function of the number and size of hash outputs **1460** rather than of key elements **1411**.

(155) The addend constructor **1424** involves producing a temporary addend value **1426** comprising a set of addend elements **1427**, labeled a.sub.0 through a.sub.x. The number x of addend elements **1427** is the same as the number of mask elements **631**, and each addend element **1427** is of the same length as a corresponding mask element **1431**. For

instance, in an embodiment where each mask element **1431** is a word, each addend element **1427** would also be a word. (156) The addend constructor **1424** produces the addend value **1426** by, for each mask element **1431**, performing a logical operation between the mask element **1431** and the folded key to produce a corresponding one of the addend elements **1427**. The logical operation may be, for instance, an AND operation or XOR operation. For example, if the chosen logical operation is AND, the addend constructor **1424** may be described as follows: $a.sub.i = XOR.sub.result \& m.sub.i$, where i is the order of the element in the corresponding set and $XOR.sub.result$ is the folded key. Thus, for instance, an AND operation between $m.sub.0$ and the folded key would produce $a.sub.0$. Or, as another example, the addend constructor may be $a.sub.i = XOR.sub.result \text{ xor } m.sub.i$.

(157) Hash outputs **1460** from a hashing component, such as hash outputs **1160a-n**, is input into hash output manipulation block **1470**. Each hash output **1460** is a different hash value computed by a separate hashing component for the hash key **1410**. There may be any number of hash outputs **1460** that are input into the hash output manipulation block **1470**, depending on the embodiment.

(158) The addend value **1426** is inputted into addition function **1428**, along with the hash output(s) **1460**. The addition function **1428** performs an addition operation between each hash output **1460** and at least a portion of the addend value **1426** to produce a transformed hash output **1480**.

(159) In an embodiment with multiple hash outputs **1460**, the entire addend **1426** may be added to each hash output **1460**, to produce each separate transformed hash output **1480**. In another embodiment with multiple hash outputs **1460**, the addend is divided into a number z portions, equal in number to the number of hash outputs **1460**. Each portion of the addend is the same size as a corresponding hash output **1460**, to which it will be added to form the corresponding transformed hash output **1480**. Hence, if each hash output **1460** is thirty-two bits, the addend may be divided into three portions, each thirty-two bits in size.

(160) From a different perspective, the hash outputs **1460** may be concatenated together to form a single value, to which the entire addend **1426** is added. The output of the addition operation may then be separated again into each transformed hash output **1480**.

(161) In an embodiment, the addition of any distinct addend portion and hash output **1460** may be such that the operation does not impact any other hash output **1460**. In this manner, the operation is said to be a modular addition operation. Other rules may also be imposed, such as a requirement that the result not saturate, and a requirement that the addition operation must simply roll over. In an embodiment where each addend element **1427** is a word, the addition operation is a word-wise addition operation.

(162) System **1400** illustrates only one of many possible arrangements of components configured to transform a hash output **1460** into a transformed hash output **1480**. Other arrangements may include fewer, additional, or different components, and the division of work between the components may vary depending on the arrangement.

(163) 1.12. Network Use Case

(164) FIG. 7 illustrates a system **700** of multiple interconnected network devices **790a-n**, collectively referred to as devices **790**, configured to utilize a same hashing function to perform a same networking function, such as assigning traffic flows or paths, but with different input and/or output manipulations, according to an embodiment. For instance, each device may utilize a hash-based lookup operation for determining paths to assign to data packets. Although each device **790** comprises the same input and/or output manipulation logic, each device **790** is configured with a different set of one or more mask values **730** (e.g. input mask values, output mask values, extension values, etc.) to use when manipulating the inputs and/or outputs of the hashing function, thus resulting in each device utilizing a different strategy to transform a given input key and/or given output hash value. Thus, for example, each device may transform a given input key differently, resulting in each device potentially determining a different hash value for the given key, even though the same hashing function is utilized at each device. Or, as another example, each device may transform a given output hash value differently, resulting in each device utilizing a different index to locate information needed for operating decisions while performing the networking function, even though the same hashing function is utilized at each device. Consequently, a device **790** may treat data packets having the same key differently when performing the operation (e.g. make different data flow decisions).

(165) 1.13. Miscellaneous

(166) According to an embodiment, a computing device may utilize hash-based operations for more than one function. For example, a network device may utilize hash-based lookup operations for more than one type of data. Hence, the device may include multiple hashing components, and multiple key transformation components. In an embodiment, some of all of these transformation components may utilize the same mask value in executing their respective transformation logic.

(167) According to some embodiments, more than one hash function may be utilized to produce a single hash output. For example, one hash function may produce one portion of the hash output, while another hash function may produce another portion of the hash output. For instance, in an embodiment, three hash functions may be utilized to produce two hash outputs. There may be two CRC-16 functions (CRC16-CCITT and CRC16-SCSI) and one CRC-32 function (CRC-32C), where each hash function receives a unique transformed key such that two 16-bit hash outputs are produced and one 32-bit hash output is produced. The two 16-bit hash outputs may be combined to yield one 32-bit hash value and the other 32-bit hash output is provided directly.

(168) Many of the techniques described herein are, for convenience, described with respect to data units referred to as data packets or network packets. However, it will be understood that the packets may be any unit of transmission over a network at any communication protocol layer, including IP packets, cells, frames, or other datagrams.

2.0. Functional Overview

(169) 2.1. Generating a Hash Value

(170) FIG. **8** illustrates an example flow **800** for generating a hash value, according to an embodiment. The various elements of flow **800** may be performed in a variety of systems, including systems such as systems **100**, **200**, **300**, **400**, and/or **500** described above. In an embodiment, each of the processes described in connection with the functional blocks described below may be implemented using one or more FPGAs, ASICs, computer programs, other software elements, and/or digital logic in any of a general-purpose computer or a special-purpose computer, while performing data retrieval, transformation, and storage operations that involve interacting with and transforming the physical state of memory of the computer.

(171) Block **810** comprises receiving an input hash key, such as key **110**. Block **820** comprises executing transformation logic based on the key and the mask value to generate a transformed key, such as transformed key **140**. The transformation logic may be any suitable logic, including the example logic described elsewhere herein. The transformation logic may involve, for instance, extension logic using an extension value **1316**, and/or logical operations with respect to an input mask value, such as predefined input mask value **130**.

(172) Block **830** comprises executing a hash function on the transformed key to generate a hash value, such as hash value **160**. Again, any suitable hash function may be utilized, as described elsewhere herein.

(173) Block **840** comprises manipulating the hash value, using hash output manipulation logic such as described elsewhere, to generate a transformed hash value. Block **840** may comprise, for instance, masking the hash value using an output mask value such as mask value **1130** and/or the key or a value derived therefrom.

(174) Block **850** comprises performing a hash-based operation using the transformed hash value, such as the lookup operation described herein, or any other suitable hash-based operation.

(175) Flow **800** illustrates only one of many possible flows for generating a hash value. Other flows may include fewer, additional, or different elements, in varying arrangements. For example, block **850** may be omitted, and the hash-based operation may be performed on the hash value directly. Or, block **820** may be omitted, and the hash function of block **830** may be performed on the input hash key directly.

(176) 2.2. Manipulating the Hash Key

(177) FIG. **9** illustrates an example flow **900** for transforming a hash key, according to an embodiment. The various elements of flow **900** may be performed in a variety of systems, including systems such as system **600** described above. In an embodiment, each of the processes described in connection with the functional blocks described below may be implemented using one or more FPGAs, ASICs, computer programs, other software elements, and/or digital logic in any of a general-purpose computer or a special-purpose computer, while performing data retrieval, transformation, and storage operations that involve interacting with and transforming the physical state of memory of the computer.

(178) Block **910** comprises inputting a hash key comprising multiple key elements, such as the key elements **611** of key **610**.

(179) Blocks **912** and **914**, which are optional depending on the embodiment, comprise identifying an extension value, such as extension value **1318**, and extending the hash key using the extension value, respectively. Block **912** may be a trivial read operation on a fixed register or other memory location, in the case that the same extension value is always used for any hash key processed by the device. Or, block **912** may comprise determining a suitable location from which to read the extension value, in embodiments where different extension values may be used to extend hash keys in different contexts (e.g. for lookups against different table). Such a determination may be made based on configuration data for the context (e.g. an identifier associated with the table).

(180) Block **920** comprises performing a folding XOR operation, such as operation **622**, to fold the key elements, including any added or modified through the optional extension in block **914**, into a folded key.

(181) Block **925** comprises identifying an input mask value, such as mask value **130**. Block **925** may be a trivial read operation on a fixed register or other memory location, in the case that the same input mask value is always used for any hash key processed by the device. Or, block **925** may comprise determining a suitable location from which to read the mask value, in embodiments where different mask values may be used to mask hash keys in different contexts (e.g. for lookups against different table). Such a determination may be made based on configuration data for the context (e.g. an identifier associated with the table).

(182) Block **930** comprises performing a multi-masking function, such as function **624**, to compute a masked addend, such as addend **626**. The multi-mask function comprises various substeps, repeated for each element of the identified mask value. At block **932**, a next mask element is identified. Block **934** comprises performing a logical operation between the next mask element and the folded key to generate a next element of the masked addend. The logical operation may be, for instance, AND or XOR, depending on the embodiment. Block **936** comprises determining whether all of the mask elements have been processed. If not, flow returns to block **932**. Otherwise, the multi-mask function is complete.

(183) Block **940** comprises performing an addition operation, such as addition function **628**, between the input key and

the masked addend, to produce a transformed key, such as transformed key **640**. The addition operation comprises various substeps, repeated for each element of a masked addend. At block **942**, a next addend element is identified. At block **944**, a next key element is identified.

(184) Block **946** comprises performing an addition operation between the addend element and the key element to produce the next element of the transformed key. As explained elsewhere herein, the addition operation may be constrained in one or more aspects. For example, the addition operation may be a bitwise addition operation and/or not permitted to saturate or roll-over. In an embodiment, the addition operation may be defined on any suitable boundary, not just bytes. For example, the operation may be defined on two-byte units, four bit units, or even units of different sizes.

(185) In an embodiment, the addition operation may involve multiple addends of potentially varying sizes. For example, an operation may involve first adding a first addend, `addend_1`, bitwise with the key elements that are also byte-wide elements. Then, or in parallel, the operation may add to the transformed key a second, two-byte addend, `addend_2`, to two bytes of the transformed key. This can be extended to any number of layers of addends.

(186) Block **948** comprises determining whether all of the key elements have been processed. If not, flow returns to block **942**. Otherwise, the addition operation is complete. The output of the addition operation is a transformed key ready for input into a hashing function.

(187) Flow **900** illustrates only one of many possible flows for generating a hash value. Other flows may include fewer, additional, or different elements, in varying arrangements. For example, while certain blocks are depicted serially, it will be recognized that a number of blocks may be performed in parallel. For instance, each iteration of blocks **932-936** or blocks **942-948** may be performed in parallel with each other. Moreover, the order of certain depicted operations may vary.

(188) 2.3. Manipulating Hash Output

(189) FIG. **15** illustrates an example flow **1500** for transforming one or more hash values, according to an embodiment. The various elements of flow **1500** may be performed in a variety of systems, including systems such as system **600** described above. In an embodiment, each of the processes described in connection with the functional blocks described below may be implemented using one or more FPGAs, ASICs, computer programs, other software elements, and/or digital logic in any of a general-purpose computer or a special-purpose computer, while performing data retrieval, transformation, and storage operations that involve interacting with and transforming the physical state of memory of the computer.

(190) Block **1510** comprises inputting a hash key comprising multiple key elements, such as the key elements **1411** of key **1410**. The hash key may be an original hash key, or an extended hash key, depending on the embodiment. Block **1520** comprises performing a folding XOR operation, such as operation **1422**, to fold the key elements into a folded key.

(191) Block **1525** comprises identifying an output mask value, such as mask value **1130**. Block **1525** may be a trivial read operation on a fixed register or other memory location, in the case that the same input mask value is always used for any hash key processed by the device. Or, block **1525** may comprise determining a suitable location from which to read the mask value, in embodiments where different mask values may be used to mask hash keys in different contexts (e.g. for lookups against different table). In an embodiment, such a determination may be made based on configuration data for the context (e.g. an identifier associated with the table). In an embodiment, the output mask value may be selected based on, for example, a given flow or set of flows maintained by a flow table or flow group table, an aggregation group, or any other identifying feature that remains fixed for packets of the same flow.

(192) Block **1530** comprises performing an addend constructor operation, such as function **1424**, to compute an addend, such as addend **1426**. The addend constructor operation comprises various substeps, repeated for each element of the identified mask value. At block **1532**, a next mask element is identified. Block **1534** comprises performing a logical operation between the next mask element and the folded key to generate a next element of the addend. The logical operation may be, for instance, AND or XOR, depending on the embodiment. Block **1536** comprises determining whether all of the mask elements have been processed. If not, flow returns to block **1532**. Otherwise, the addend constructor operation is complete.

(193) Block **1538** comprises receiving one or more hash values as input. The one or more hash values are hash values that have been calculated based on the key that was inputted in block **1510**. The one or more hash values are different values. For instance, the one or more hash values may have been calculated using different hash functions.

(194) Block **1540** comprises performing an addition operation, such as addition function **1428**, between the one or more hash values and the addend, to produce a transformed hash value, such as transformed hash value **1180**. The addition operation comprises various substeps, repeated for each of the hash values. At block **1542**, a next addend portion is identified. Depending on the embodiment, the portion may be the entire addend, or a number of bits equal to the size of the next hash value to process, or other portions of the addend. At block **1544**, a next hash value to transform is identified. Block **1546** comprises performing an addition operation between the addend element and the hash value to produce the next transformed hash value. As explained elsewhere herein, the addition operation may be constrained in one or more aspects. For example, the addition operation may be a modular addition operation and/or not permitted to saturate or roll-over. Again, the addition operation may be defined on boundaries of any suitable size, and/or involve

multiple adds. Block **1548** comprises determining whether all of the hash values have been processed. If not, flow returns to block **1542**. Otherwise, the addition operation is complete.

(195) At block **1550**, the result of the addition operation is outputted. The output of the addition operation is a set of one or more transformed hash values ready for input into a hash-based operation.

(196) Flow **1500** illustrates only one of many possible flows for generating a hash value. Other flows may include fewer, additional, or different elements, in varying arrangements. For example, while certain blocks are depicted serially, it will be recognized that a number of blocks may be performed in parallel. For instance, each iteration of blocks **1532-1536** or blocks **1542-1548** may be performed in parallel with each other. Moreover, the order of certain depicted operations may vary.

(197) 2.4. Using Transformed Keys for Lookup Operations in a Network Device

(198) FIG. **10** illustrates an example flow **1000** for using transformed keys to support example lookup operations in a network device, according to an embodiment. The various elements of flow **1000** may be performed in a variety of systems, including systems such as described above. In an embodiment, each of the processes described in connection with the functional blocks described below may be implemented using one or more FPGAs, ASICS, computer programs, other software elements, and/or digital logic in any of a general-purpose computer or a special-purpose computer, while performing data retrieval, transformation, and storage operations that involve interacting with and transforming the physical state of memory of the computer.

(199) Block **1005** comprises starting or initializing the network device. For example, the device may be powered on after being in a non-powered state, or an operating system of the device may be rebooted. Block **1010** comprises, during said starting or initializing, calculating and storing a random mask value for the device. This mask value may serve as a predefined mask value (e.g. mask value **130**) for one or more key transformations within the device.

(200) Block **1015** comprises receiving a data packet, such as an Ethernet packet, IP packet, or other datagram. For example, such a data packet may be received over a communication interface of the network device from another network device on the same network. The network device may be configured to route the data packet to yet another device on the network. Routing logic within the network device may be configured to utilize one or more hashing operations when determining where to send the data packet, when to send the data packet, whether to manipulate the data packet, or otherwise handle the data packet.

(201) Block **1020** comprises receiving a request to lookup information associated with the data packet. Such lookup requests may occur in a variety of contexts, such as attempts to determine a flow or path to assign the data packet to. The information is information stored in an entry within one or more lookup tables, such as tables **270**, configured to store a specific type of information, such as memory addresses, assigned traffic flows, or other suitable data.

(202) Block **1025** comprises determining a key for the data packet. The key may be generated, for instance, by taking a first number of bits from the data packet, extracting data from specific locations within the packet header or payload, such as source and destination fields, applying some function to the packet, or any other suitable means.

(203) Block **1030** comprises transforming the key using transformation logic based on the mask value, such as described elsewhere herein. Block **1035** comprises computing multiple hash values using the transformed key, by applying different hash functions to the transformed key. Block **1040** comprises, for each hash value, searching for the generated key within the table slot whose index corresponds to hash value. Block **1045** comprises, once the key is found, returning a value stored in association with the generated key. Alternatively, if the key is not found, block **1050** comprises returning data indicating that no value is stored for the key.

(204) Block **1055** comprises making a routing decision for the data packet based on the returned value (or lack thereof).

(205) Flow **1000** illustrates only one of many possible flows for using transformed keys to support lookup operations in a network device. Other flows may include fewer, additional, or different elements, in varying arrangements. For example, in an embodiment, instead of block **1050**, flow **1000** may comprise steps such as, when the key is not found, calculating a value for the key and writing the key and the associated value at an index corresponding to one of the computed hash values. The calculated value may then be returned. Such steps may further comprise moving previously stored keys to other indexes corresponding to other hash values calculated for those keys in order to accommodate the newly written key, when needed.

(206) 2.5. Using Hash Manipulations for Mapping Functions

(207) FIG. **16** illustrates an example flow **1600** for using transformed keys to support example mapping operations in a network device, according to an embodiment. The various elements of flow **1600** may be performed in a variety of systems, including systems such as described above. In an embodiment, each of the processes described in connection with the functional blocks described below may be implemented using one or more FPGAs, ASICS, computer programs, other software elements, and/or digital logic in any of a general-purpose computer or a special-purpose computer, while performing data retrieval, transformation, and storage operations that involve interacting with and transforming the physical state of memory of the computer.

(208) Blocks **1605-1625** proceed as with blocks **1005-1025**, except that the request of block **1620** is to identify a value mapped to the data packet in a specific mapping table (e.g. in a table of assigned paths or traffic flows).

(209) Block **1628** comprises extending the key using an extension value, as described in other sections.

(210) Block **1630** comprises transforming the extended key using an input mask value, as described in other sections.

Block **1635** comprises computing multiple hash values using the transformed key. For example, in an embodiment, three different outputs are generated using three different hashing functions.

(211) Block **1640** comprises transforming the generated hash values using the key, extended key, or transformed key along with an output mask value. Examples of logic for transforming the generated hash values are described in other sections.

(212) Block **1645** comprises selecting which of the multiple transformed hash values to use for the mapping table, typically using configuration data associated with the mapping table. Generally, the same transformed hash output should be selected whenever mapping values from the mapping table.

(213) Block **1650** comprises reading a mapped value from the mapping table, at an index indicated by the hash value. The index may be equal to the transformed hash value itself, or, if the transformed hash value is greater than the number of buckets within the mapping table, the index may be the remainder produced by a modulo operation dividing the transformed hash value by the number of buckets.

(214) Block **1655** comprises returning the mapped value.

(215) Flow **1600** illustrates only one of many possible flows for using transformed keys to support mapping operations in a network device. Other flows may include fewer, additional, or different elements, in varying arrangements. For example, in an embodiment, instead of storing an actual mapping table, the mappings may be represented by other suitable data structures. For example, a data structure may map ranges of index values to mapped values, and the mapping may comprise determining which range includes the determined index.

3.0. Example Embodiments

(216) Examples of some embodiments are represented, without limitation, in the following clauses:

(217) According to an embodiment, an apparatus comprises a transformation component, configured to input a hash key, generate a transformed hash key based on the inputted hash key, and output the transformed hash key; and a hashing component, configured to input the transformed hash key, generate a hash value based on applying a hash function to the transformed hash key, and output the hash value.

(218) According to an embodiment, an apparatus comprises a hashing component, configured to input a hash key, generate a hash value based on applying a hash function to the transformed hash key, and output the hash value; and a hash value output manipulation component, configured to input a hash value, generate a transformed hash value based on the inputted hash value, and output the transformed hash value.

(219) According to an embodiment, an apparatus comprises: key generation logic configured to identify keys for data units; a hash block configured to input the keys and output indexes, the hash block further configured to, for a given key of the keys, determine one or more of the indexes that hash to the given key using one or more hashing functions applied to one or more transformed keys generated from the given key and one or more selected pre-defined mask values; one or more lookup components configured to search for values associated with the keys in different sets of one or more lookup tables, the one or more lookup components searching for a given key in a given set of one or more lookup tables at one or more of the indexes outputted by the hash block for the given key.

(220) According to an embodiment, an apparatus comprises: key generation logic configured to identify keys for data units; a hash block configured to input the keys and output indexes, the hash block further configured to, for a given key of the keys, determine one or more of the indexes that hash to the given key using: one or more hashing functions applied to the key or one or more key values derived from the given key, and hash value manipulation logic that applies a selected pre-defined output mask value to one or more hash values output by the one or more hashing functions; one or more mapping components configured to search for values that map to the keys in mapping tables, the one or more mapping components locating the value mapped to a given key in a given mapping table at one or more of the indexes outputted by the hash block for the given key.

(221) According to an embodiment, an apparatus comprises: a transformation component, configured to input a hash key comprising one or more key elements, generate a transformed hash key based on a series of logical operations involving the one or more key elements of the inputted hash key and a fixed mask value comprising a corresponding one or more mask elements, and output the transformed hash key; a hashing component, configured to input the transformed hash key, generate a hash value based on applying a hash function to the transformed hash key, and output the hash value.

(222) In an embodiment, the one or more key elements include two or more key elements. In an embodiment, the one or more key elements are each of the same length as a corresponding mask element of the one or more mask elements. In an embodiment, each of the one or more key elements is of an equal length. In an embodiment, each of the one or more key elements is a byte. In an embodiment, wherein the hash key is different than the transformed hash key, and the fixed mask value is different than the hash key and the transformed hash key.

(223) In an embodiment, the transformation component and the hashing component are implemented by one or more FPGAs or ASICs.

(224) In an embodiment, the series of logical operations includes an addition operation between the hash key and an addend determined based on the fixed mask value. In an embodiment, the series of logical operations includes an XOR-folding operation between each of the one or more key elements to produce a folded key value, AND or XOR operations between the folded key value and each mask element of the one or more mask elements to produce a masked

addend, and an addition operation between the masked addend and the hash key to produce the transformed hash key.

(225) In an embodiment, an apparatus further comprises: a mask value determination component, configured to initialize the fixed mask value upon a restart operation of the apparatus. In an embodiment, the mask value determination component is configured to initialize the fixed mask value to a random value upon the restart operation. In an embodiment, the fixed mask value does not change between system restart operations.

(226) In an embodiment, the hash function is a cyclic redundancy check (CRC) function.

(227) In an embodiment, an apparatus further comprises: one or more lookup tables comprising entries that store values associated with transformed hash keys; a lookup component configured to use the hash value generated by the hashing component to identify an index of an entry in the one or more lookup tables in which a value associated with the hash key is stored.

(228) In an embodiment, the hash key is derived from a data packet and the value is a memory address in which information related to the data packet is stored. In an embodiment, the information related to the data packet is an assigned path or flow. In an embodiment, the information related to the data packet is used for packet forwarding actions, including one or more of: signature identification, counter identification, path resolution, group classification, security, or editing.

(229) In an embodiment, an apparatus further comprises: a key generation component configured to input a data packet, generate the hash key based on the data packet, and output the hash key. In an embodiment, the key generation component is configured to generate the hash key by using one or more fields of the inputted data packet as the one or more hash key elements. In an embodiment, the one or more fields include one or more of: a destination field, a source field, a protocol field, a destination port field, or a source port field.

(230) In an embodiment, an apparatus further comprises: an extension component configured to prepend and/or append one or more fixed extension values to the hash key prior to the hash key being inputted to the transformation component.

(231) In an embodiment, the hashing component is a first hashing component of multiple hashing components configured to input the transformed hash key, each of the multiple hashing components configured to output a different hash value generated based upon the transformed hash key, using a different hash function of the transformed hash key.

(232) In an embodiment, the transformation component is a first transformation component of multiple transformation components configured to input the hash key, each of the multiple transformation components configured to generate a different transformed hash key based upon the hash key, using a different fixed mask value.

(233) In an embodiment, the hashing component is a first hashing component of multiple hashing components, each of the multiple hashing components configured to input a different one of the different transformed hash keys and output a different hash value generated based upon the respectively inputted transformed hash key.

(234) In an embodiment, the hashing component is a first hashing component of multiple hashing components, including a second hashing component configured to output another hash value generated based on applying a hash function to the hash key.

(235) In an embodiment, an apparatus further comprises: a lookup component configured to utilize each of the different hash values outputted by the multiple hashing components as indexes in one or more lookup tables, the lookup component searching for an entry associated with the hash key at each of the indexes.

(236) In an embodiment, the lookup component is configured to, if no entry associated with the hash key is found at any of the indexes, add an entry associated with the hash key at one of the indexes. In an embodiment, the lookup component is configured to, if no entry associated with the transformed hash key is found at any of the indexes, and the lookup table is full at each of the indexes, move an entry associated with an other hash key from a particular one of the indexes to another index to which one of the multiple hashing components hashes an other transformed hash key generated by the transformation component for the other hash key, and add an entry associated with the hash key at the particular one of the indexes.

(237) In an embodiment, an apparatus further comprises a plurality of sets of lookup tables, each set of lookup tables having configuration data indicating one or more of: which of the different outputted hash values the lookup component is to utilize as the index when searching in the set of lookup tables, how many buckets are in the set of lookup tables, or the mask value to use to compute the transformed hash key when searching in the set of lookup tables.

(238) In an embodiment, an apparatus further comprises a plurality of sets of lookup tables, each set of lookup tables having configuration data indicating one or more of: which hash function is to be applied to generate the hash value when searching in the set of lookup tables, or a manner by which the hash value is to be manipulated.

(239) In an embodiment, an apparatus further comprises a plurality of lookup components, each lookup component configured to utilize at least one of the different hash values generated by the multiple hashing components as an index in a lookup table corresponding to the lookup component, each lookup component having configuration data indicating which of the different outputted hash values the lookup component is to utilize as the index.

(240) In an embodiment, an apparatus is a network device configured to switch or route Ethernet or Internet Protocol (IP) packets, the network device further comprising: a key generation component configured to generate hash keys to input into the transformation component based on the packets; a lookup component configured to locate information associated with the packets using hash values generated by the hashing component; packet handling logic configured to

determine how to handle the packets based on the information located by the lookup component.

(241) According to an embodiment, a system comprises: at least a plurality of network devices interconnected by one or more networks, each device in the plurality comprising: one or more network interfaces configured to receive and send messages to other devices over the one or more networks; key generation logic configured to generate hash keys based on the messages; a transformation component, configured to input hash keys, generate transformed hash keys based on the inputted hash keys, and to output the transformed hash keys; a hashing component, configured to input the transformed hash keys, generate hash values based on applying a hash function to the transformed hash key, and output the hash values; and message handling logic configured to handle the messages based on the hash values; wherein the transformation component on each device is configured to apply a different transformation to the hash keys used to generate the hash values for the message handling logic at the device than applied at each other of the devices in the plurality.

(242) In an embodiment, the transformation component at each device is configured to utilize a mask value to transform the hash keys generated at the device, each device utilizing a different mask value to transform the hash keys used to generate the hash values for the message handling logic at the device.

(243) In an embodiment, the transformation component transforms hash keys by, for each given hash key of the hash keys, generating an folded key using XOR operations between each of one or more key elements within the given hash key, generating a masked addend using AND or XOR operations between the folded key and each mask element of one or more mask elements in a fixed mask value associated with the transformation component, generating a given transformed hash key using an addition operation between the masked addend and the given hash key.

(244) In an embodiment, the different mask value of each device is randomly assigned when the device is booted.

(245) In an embodiment, the messages are IP packets or Ethernet packets, wherein the hash keys are predefined locations within the IP packets or Ethernet packets. In an embodiment, the predefined locations include one or more of a source address field, source port field, destination address field, destination source field, or protocol type field.

(246) According to an embodiment, an apparatus comprises: a key generation configured to generate a hash key from a data unit; a key extension component configured to extend the hash key by prepending or appending a predefined extension value to the hash key; a key transformation component configured to transform the extended hash key based on an input mask value; a hashing component configured to generate one or more hash values by performing one or more hash functions on the transformed extended hash key; a hash value manipulation component configured to transform the one or more hash values based on an output mask value and at least one of the extended hash key or the transformed extended hash key; a hash-based operation component configured to perform a hash-based operation with respect to the data unit based on information stored at a location whose address is determined using a particular transformed hash value generated by the hash value manipulation component.

(247) In an embodiment, the data unit is a network packet being routed through the apparatus, wherein the hash-based operation includes identifying a path or flow to assign to the network packet in a mapping table, the assigned path or flow stored at an entry within the mapping table whose index corresponds to the particular transformed hash value.

(248) In an embodiment, the key generation component is configured to generate the hash key by extracting at least one or more of a source address, destination address, source port, destination port, or protocol of the network packet.

(249) In an embodiment, the hashing component is configured to generate two or more hash values based on the transformed extended hash key, the hash value manipulation component is configured to transform the two or more hash values, and the hash-based operation component is configured to select which of the two or more hash values to input as the particular transformed hash value based on configuration data associated with a mapping table from which the information is retrieved.

(250) In an embodiment, the extension value, input mask value, and output mask value are generated at least partially randomly whenever the apparatus boots or restarts.

(251) In an embodiment, an apparatus further comprises logic configured to determine where to read the extension value, input mask value, and output mask value from based on configuration data associated with a mapping table from which the information is retrieved and/or a group to which the data unit is assigned.

(252) In an embodiment, the hash value manipulation component is configured to transform the hash value by generating a first folded key value from the extended hash key, the first folded key value generated by dividing the inputted extended hash key into a plurality of first key portions and performing a first logical operation between each of the first key portions, generating an output addend based on second logical operations between the first folded key and each of two or more mask elements of the output mask value, and generating a transformed hash value based on a first addition operation between at least a portion of the output addend and the hash value.

(253) In an embodiment, the hashing component is configured to generate two or more hash values based on the transformed extended hash key, and the hash value manipulation component is configured to transform the two or more hash values by adding each of two or more distinct sub-elements of the output addend to a corresponding one of the two or more hash values.

(254) In an embodiment, the hash-based operation component is configured to determine which of the multiple transformed hash values to input based upon configuration data associated with a table for the hash-based operation component is instructed to use an output of the hash output manipulation component as an index during the hash-based

operation.

(255) In an embodiment, the extended hash key is dividable into a plurality of second key portions, and the key transformation component is configured to: perform a third logical operation between each of the second key portions to produce a second folded key value, generate a input addend based on fourth logical operations between the second folded key and each of two or more mask elements of the input mask value, and generate the transformed extended hash key based on a second addition operation between the input addend and the extended hash key.

(256) In an embodiment, the first logical operation and the third logical operation are both an XOR-folding operation, the second logical operations and the fourth logical operations are each an AND or an XOR operation, and the first addition operation and second addition operation are modular addition operations.

(257) In an embodiment, the second key portions are bytes and the first key portions are words.

(258) In an embodiment, the one or more hashing functions include at least one hardware-implemented cyclic-redundancy-check (CRC) based function.

(259) In an embodiment, the key generation, key extension component, key transformation component, hashing component, and hash value manipulation component are implemented by one or more FPGAs or ASICs.

(260) According to an embodiment, a system comprises: at least a plurality of network devices interconnected by one or more networks, each device in the plurality comprising: one or more network interfaces configured to receive and send messages to other devices over the one or more networks; key generation logic configured to generate hash keys based on the messages; a hashing component, configured to generate hash values based on applying a hash function to either the hash keys or manipulated hash keys derived from the hash keys, and to output the hash values; a hash value manipulation component configured to input the hash values, generate transformed hash values based on the hash values, and output the transformed hash values; and message handling logic configured to handle the messages based on the transformed hash values; wherein the hash value manipulation component on each device is configured to apply a different transformation to the hash values for the message handling logic at the device than applied at each other of the devices in the plurality.

(261) In an embodiment, the hash value manipulation component at each device is configured to utilize a mask value to transform the hash values generated at the device, each device utilizing a different mask value to transform the hash values for the message handling logic at the device.

(262) In an embodiment, the transformation component transforms hash values by, for each given hash key of the hash keys: generating an folded key using XOR operations between each of two or more key elements within the given hash key, generating a masked addend using AND or XOR operations between the folded key and each mask element of two or more mask elements in a fixed mask value associated with the hash value manipulation component, and generating a given transformed hash value using an addition operation between at least a portion the masked addend and a given hash value computed for the given hash key.

(263) In an embodiment, the different mask value of each device is randomly assigned when the device is booted.

(264) In an embodiment, the messages are IP or Ethernet packets, wherein the hash keys are extracted from predefined locations within the packets. In an embodiment, the predefined fields include two or more of a source address field, source port field, destination address field, destination source field, or protocol type field.

(265) In an embodiment, each device further comprises a hash key manipulation component, configured to input hash keys, extend the hash keys, generate transformed extended hash keys based on the inputted hash keys, and to output the transformed extended hash keys as manipulated hash keys for input into the hashing component.

(266) According to an embodiment, an apparatus comprises: a hashing component, configured to input a key value derived from a hash key, generate a hash value based on applying a hash function to the input, and output the hash value; a hash output manipulation component, configured to input the hash key and the hash value from the hashing component, generate a transformed hash value based upon the inputted hash key, a fixed mask value, and the hash value, and output the transformed hash value.

(267) In an embodiment, the hash output manipulation component and the hashing component are implemented by one or more FPGAs or ASICs.

(268) In an embodiment, the hash value is different than the transformed hash value, the hash key is different than the hash value and the transformed hash value, and the fixed mask value is different than the hash key, hash value, and the transformed hash value.

(269) In an embodiment, the hash value manipulation component is configured to generate the transformed hash value by generating an addend using the inputted hash key and a fixed mask value, and generate a transformed hash value based on an operation between the addend and the hash value.

(270) In an embodiment, the hashing component is configured to generate multiple hash values for the hash key using multiple different hash functions, wherein the hash output manipulation component is configured to input the multiple hash values, wherein the operation between the addend and the transformed hash value comprises adding different distinct addend sub-elements to different hash values of the multiple hash values to generate multiple transformed hash values, the hash output manipulation component configured to output the multiple transformed hash values.

(271) In an embodiment, the hashing component is configured to generate multiple hash values for the hash key using multiple different hash functions, wherein the hash output manipulation component is configured to input the multiple

hash values, wherein the hash output manipulation component configured to generate multiple transformed hash values based on the multiple different hash functions, the fixed mask value, and the hash key, wherein the hash output manipulation component is further configured to output the multiple transformed hash values. The apparatus further comprises: a hash-based operation component configured to input a particular hash value and perform a hash-based operation based on the inputted particular hash-based operation, wherein the hash-based operation component is configured to determine which of the multiple hash values to input from the hash output manipulation component.

(272) In an embodiment, the hash-based operation component is configured to determine which of the multiple transformed hash values to input based upon configuration data associated with a table for the hash-based operation component is instructed to use an output of the hash output manipulation component as an index during the hash-based operation.

(273) In an embodiment, the hash-based operation component is configured to determine which hash function is to be applied to generate the hash value, or a manner by which the hash value is to be manipulated, based upon configuration data associated with a table for the hash-based operation component is instructed to use an output of the hash output manipulation component as an index during the hash-based operation.

(274) In an embodiment, the hash-based operation component is configured to determine which of the multiple transformed hash values to input based upon configuration data associated a characteristic of a data unit for which the hash key was generated, the hash-based operation being performed with respect to the data unit.

(275) In an embodiment, an apparatus further comprises a key transformation component configured to input the hash value, generate the key value based on transforming the hash value, and output the key value to the hashing component.

(276) In an embodiment, the hash key is an extended hash key, and the apparatus further comprises: a key extension component configured to input an original key, extend the original key by prepending or appending a fixed extension value to the original key, and output the hash key.

(277) In an embodiment, the hash value manipulation component is configured to generate the transformed hash value by generating a folded key by performing a first logical operation between each of two or more key elements of the inputted hash key, generating an addend based on second logical operations between the folded key and each of two or more mask elements of a fixed mask value, and generating a transformed hash value based on an operation between at least a portion of the addend and the hash value.

(278) In an embodiment, the two or more key elements and two or more mask elements are each of the same length. In an embodiment, each of the two or more key elements and two or more mask elements are sixteen bits in length. In an embodiment, the operation between the addend and the hash value is an addition operation. In an embodiment, the first logical operation is an XOR-folding operation, the second logical operations are each an AND or an XOR operation.

(279) In an embodiment, an apparatus further comprises a mask value determination component, configured to initialize the fixed mask value upon a restart operation of the apparatus.

(280) In an embodiment, the mask value determination component is configured to initialize the fixed mask value to a random value upon the restart operation. In an embodiment, the fixed mask value does not change between system restart operations.

(281) In an embodiment, the hash function is a cyclic redundancy check (CRC) function.

(282) In an embodiment, an apparatus further comprises one or more mapping tables comprising entries that store values associated with transformed hash keys;

(283) a mapping component configured to use the transformed hash value generated by the hashing component to identify an index of an entry in the one or more lookup tables in which a value associated with the hash key is stored.

(284) In an embodiment, the hash key is derived from a data packet and the value is a memory address in which information related to the data packet is stored. In an embodiment, the information related to the data packet is an assigned path or flow. In an embodiment, the information related to the data packet is used for packet forwarding actions, including one or more of: signature identification, counter identification, path resolution, group classification, security, or editing.

(285) In an embodiment, an apparatus further comprises a key generation component configured to input a data packet, generate the hash key based on the data packet, and output the hash key.

(286) In an embodiment, the key generation component is configured to generate the hash key by using two or more fields of the inputted data packet as the two or more hash key elements. In an embodiment, the two or more fields include two or more of: a destination field, a source field, a protocol field, a destination port field, or a source port field.

(287) In an embodiment, an apparatus further comprises a manipulation component configured to manipulate the hash key prior to the hash key being inputted to the hashing component, the manipulating including one or more of prepending and/or appending one or more fixed extension values to the hash key, or performing a series of logical operations based on the hash key.

(288) In an embodiment, the hashing component outputs a plurality of different hash values generated based upon the hash key, using a different hash function.

(289) In an embodiment, the apparatus is a network device configured to switch or route Ethernet or Internet Protocol (IP) packets, the network device further comprising: a key generation component configured to generate hash keys to input into the transformation component based on the packets; a mapping component configured to locate information

associated with the packets using hash values generated by the hashing component; packet handling logic configured to determine how to handle the packets based on the information located by the mapping component.

(290) According to an embodiment, a system comprises: at least a plurality of network devices interconnected by one or more networks, each device in the plurality comprising: one or more network interfaces configured to receive and send messages to other devices over the one or more networks; key generation logic configured to generate hash keys based on the messages; a hashing component, configured to input the hash keys, generate hash values based on applying a hash function to the hash key, and output the hash values; and a transformation component, configured to input hash values, generate transformed hash values based on the inputted hash values, and to output the transformed hash values; message handling logic configured to handle the messages based on the transformed hash values; wherein the transformation component on each device is configured to apply a different transformation to the hash values used to generate the hash values for the message handling logic at the device than applied at each other of the devices in the plurality.

(291) In an embodiment, the transformation component at each device is configured to utilize a mask value to transform the hash keys generated at the device, each device utilizing a different mask value to transform the hash keys used to generate the hash values for the message handling logic at the device.

(292) In an embodiment, the different mask value of each device is randomly assigned when the device is booted.

(293) In an embodiment, the messages are IP packets or Ethernet packets, wherein the hash keys are predefined locations within the IP packets or Ethernet packets.

(294) According to an embodiment, a method comprises: receiving input identifying a hash key comprising one or more key elements; generating a transformed hash key based on a series of logical operations involving the one or more key elements of the inputted hash key and a fixed mask value comprising a corresponding one or more mask elements; generating a hash value based on applying a hash function to the transformed hash key; performing a lookup operation for information associated with the hash key using the generated hash value.

(295) According to an embodiment, a method comprises: generating a hash key from a data unit; extending the hash key by prepending or appending a predefined extension value to the hash key; transforming the extended hash key based on an input mask value; generating one or more hash values by performing one or more hash functions on the transformed extended hash key; transforming the one or more hash values based on an output mask value and at least one of the extended hash key or the transformed extended hash key; performing a hash-based operation with respect to the data unit based on information stored at a location whose address is determined using a particular transformed hash value generated by the hash value manipulation component.

(296) According to an embodiment, a method comprises: inputting a key value derived from a hash key; generating a hash value based on applying a hash function to the inputted key value; generating a transformed hash value based upon the inputted hash key, a fixed mask value, and the hash value; searching for information mapped to the hash key using the transformed hash value.

(297) Other examples of these and other embodiments are found throughout this disclosure.

4.0. Implementation Mechanism-Hardware Overview

(298) According to one embodiment, the techniques described herein are implemented by one or more special-purpose computing devices. The special-purpose computing devices may be desktop computer systems, portable computer systems, handheld devices, networking devices, or any other device that incorporates hard-wired and/or program logic to implement the techniques. The special-purpose computing devices may be hard-wired to perform the techniques, or may include digital electronic devices such as one or more application-specific integrated circuits (ASICs) or field programmable gate arrays (FPGAs) that are persistently programmed to perform the techniques. Such special-purpose computing devices may also combine custom hard-wired logic, ASICs, or FPGAs with custom programming to accomplish the techniques.

(299) Though the foregoing techniques are described with respect to a hardware implementation, which provides a number of advantages in certain embodiments, it will also be recognized that, in another embodiment, the foregoing techniques may still provide certain advantages when performed partially or wholly in software. Accordingly, in such an embodiment, a suitable implementing apparatus comprises a general-purpose hardware processor and is configured to perform any of the foregoing methods by executing program instructions in firmware, memory, other storage, or a combination thereof.

(300) FIG. 11 is a block diagram that illustrates a computer system 1100 that may be utilized in implementing the above-described techniques, according to an embodiment. Computer system 1100 may be, for example, a desktop computing device, laptop computing device, tablet, smartphone, server appliance, computing mainframe, multimedia device, handheld device, networking apparatus, or any other suitable device.

(301) Computer system 1100 may include one or more ASICs, FPGAs, or other specialized circuitry 1103 for implementing program logic as described herein. For example, circuitry 1103 may include fixed and/or configurable hardware logic blocks for implementing some or all of the described techniques, input/output (I/O) blocks, hardware registers or other embedded memory resources such as random access memory (RAM) for storing various data, and so forth. The logic blocks may include, for example, arrangements of logic gates, flip-flops, multiplexers, and so forth, configured to generate an output signals based on logic operations performed on input signals.

(302) Additionally, and/or instead, computer system **1100** may include one or more hardware processors **1104** configured to execute software-based instructions. Computer system **1100** may also include one or more busses **1102** or other communication mechanism for communicating information. Busses **1102** may include various internal and/or external components, including, without limitation, internal processor or memory busses, a Serial ATA bus, a PCI Express bus, a Universal Serial Bus, a HyperTransport bus, an Infiniband bus, and/or any other suitable wired or wireless communication channel.

(303) Computer system **1100** also includes one or more memories **1106**, such as a RAM, hardware registers, or other dynamic or volatile storage device for storing data units to be processed by the one or more ASICs, FPGAs, or other specialized circuitry **1103**. Memory **1106** may also or instead be used for storing information and instructions to be executed by processor **1104**. Memory **1106** may be directly connected or embedded within circuitry **1103** or a processor **1104**. Or, memory **1106** may be coupled to and accessed via bus **1102**. Memory **1106** also may be used for storing temporary variables, data units describing rules or policies, or other intermediate information during execution of program logic or instructions.

(304) Computer system **1100** further includes one or more read only memories (ROM) **1108** or other static storage devices coupled to bus **1102** for storing static information and instructions for processor **1104**. One or more storage devices **1110**, such as a solid-state drive (SSD), magnetic disk, optical disk, or other suitable non-volatile storage device, may optionally be provided and coupled to bus **1102** for storing information and instructions.

(305) A computer system **1100** may also include, in an embodiment, one or more communication interfaces **1118** coupled to bus **1102**. A communication interface **1118** provides a data communication coupling, typically two-way, to a network link **1120** that is connected to a local network **1122**. For example, a communication interface **1118** may be an integrated services digital network (ISDN) card, cable modem, satellite modem, or a modem to provide a data communication connection to a corresponding type of telephone line. As another example, the one or more communication interfaces **1118** may include a local area network (LAN) card to provide a data communication connection to a compatible LAN. As yet another example, the one or more communication interfaces **1118** may include a wireless network interface controller, such as a 1102.11-based controller, Bluetooth controller, Long Term Evolution (LTE) modem, and/or other types of wireless interfaces. In any such implementation, communication interface **1118** sends and receives electrical, electromagnetic, or optical signals that carry digital data streams representing various types of information.

(306) Network link **1120** typically provides data communication through one or more networks to other data devices. For example, network link **1120** may provide a connection through local network **1122** to a host computer **1124** or to data equipment operated by a Service Provider **1126**. Service Provider **1126**, which may for example be an Internet Service Provider (ISP), in turn provides data communication services through a wide area network, such as the world wide packet data communication network now commonly referred to as the “Internet” **1128**. Local network **1122** and Internet **1128** both use electrical, electromagnetic or optical signals that carry digital data streams. The signals through the various networks and the signals on network link **1120** and through communication interface **1118**, which carry the digital data to and from computer system **1100**, are example forms of transmission media.

(307) In an embodiment, computer system **1100** can send messages and receive data through the network(s), network link **1120**, and communication interface **1118**. In some embodiments, this data may be data units that the computer system **1100** has been asked to process and, if necessary, redirect to other computer systems via a suitable network link **1120**. In other embodiments, this data may be instructions for implementing various processes related to the described techniques. For instance, in the Internet example, a server **1130** might transmit a requested code for an application program through Internet **1128**, ISP **1126**, local network **1122** and communication interface **1118**. The received code may be executed by processor **1104** as it is received, and/or stored in storage device **1110**, or other non-volatile storage for later execution. As another example, information received via a network link **1120** may be interpreted and/or processed by a software component of the computer system **1100**, such as a web browser, application, or server, which in turn issues instructions based thereon to a processor **1104**, possibly via an operating system and/or other intermediate layers of software components.

(308) Computer system **1100** may optionally be coupled via bus **1102** to one or more displays **1112** for presenting information to a computer user. For instance, computer system **1100** may be connected via an High-Definition Multimedia Interface (HDMI) cable or other suitable cabling to a Liquid Crystal Display (LCD) monitor, and/or via a wireless connection such as peer-to-peer Wi-Fi Direct connection to a Light-Emitting Diode (LED) television. Other examples of suitable types of displays **1112** may include, without limitation, plasma display devices, projectors, cathode ray tube (CRT) monitors, electronic paper, virtual reality headsets, braille terminal, and/or any other suitable device for outputting information to a computer user. In an embodiment, any suitable type of output device, such as, for instance, an audio speaker or printer, may be utilized instead of a display **1112**.

(309) One or more input devices **1114** are optionally coupled to bus **1102** for communicating information and command selections to processor **1104**. One example of an input device **1114** is a keyboard, including alphanumeric and other keys. Another type of user input device **1114** is cursor control **1116**, such as a mouse, a trackball, or cursor direction keys for communicating direction information and command selections to processor **1104** and for controlling cursor movement on display **1112**. This input device typically has two degrees of freedom in two axes, a first axis (e.g., x) and

a second axis (e.g., y), that allows the device to specify positions in a plane. Yet other examples of suitable input devices **1114** include a touch-screen panel affixed to a display **1112**, cameras, microphones, accelerometers, motion detectors, and/or other sensors. In an embodiment, a network-based input device **1114** may be utilized. In such an embodiment, user input and/or other information or commands may be relayed via routers and/or switches on a Local Area Network (LAN) or other suitable shared network, or via a peer-to-peer network, from the input device **1114** to a network link **1120** on the computer system **1100**.

(310) As discussed, computer system **1100** may implement techniques described herein using customized hard-wired logic, one or more ASICs or FPGAs **1103**, firmware and/or program logic, which in combination with the computer system causes or programs computer system **1100** to be a special-purpose machine. According to one embodiment, however, the techniques herein are performed by computer system **1100** in response to processor **1104** executing one or more sequences of one or more instructions contained in main memory **1106**. Such instructions may be read into main memory **1106** from another storage medium, such as storage device **1110**. Execution of the sequences of instructions contained in main memory **1106** causes processor **1104** to perform the process steps described herein.

(311) The term “storage media” as used herein refers to any non-transitory media that store data and/or instructions that cause a machine to operate in a specific fashion. Such storage media may comprise non-volatile media and/or volatile media. Non-volatile media includes, for example, optical or magnetic disks, such as storage device **1110**. Volatile media includes dynamic memory, such as main memory **1106**. Common forms of storage media include, for example, a floppy disk, a flexible disk, hard disk, solid state drive, magnetic tape, or any other magnetic data storage medium, a CD-ROM, any other optical data storage medium, any physical medium with patterns of holes, a RAM, a PROM, and EPROM, a FLASH-EPROM, NVRAM, any other memory chip or cartridge.

(312) Storage media is distinct from but may be used in conjunction with transmission media. Transmission media participates in transferring information between storage media. For example, transmission media includes coaxial cables, copper wire and fiber optics, including the wires that comprise bus **1102**. Transmission media can also take the form of acoustic or light waves, such as those generated during radio-wave and infra-red data communications.

(313) Various forms of media may be involved in carrying one or more sequences of one or more instructions to processor **1104** for execution. For example, the instructions may initially be carried on a magnetic disk or solid state drive of a remote computer. The remote computer can load the instructions into its dynamic memory and use a modem to send the instructions over a network, such as a cable network or cellular network, as modulated signals. A modem local to computer system **1100** can receive the data on the network and demodulate the signal to decode the transmitted instructions. Appropriate circuitry can then place the data on bus **1102**. Bus **1102** carries the data to main memory **1106**, from which processor **1104** retrieves and executes the instructions. The instructions received by main memory **1106** may optionally be stored on storage device **1110** either before or after execution by processor **1104**.

5.0. Extensions and Alternatives

(314) As used herein, the terms “first,” “second,” “certain,” and “particular” are used as naming conventions to distinguish queries, plans, representations, steps, objects, devices, or other items from each other, so that these items may be referenced after they have been introduced. Unless otherwise specified herein, the use of these terms does not imply an ordering, timing, or any other characteristic of the referenced items.

(315) In the drawings, the various components are depicted as being communicatively coupled to various other components by arrows. These arrows illustrate only certain examples of information flows between the components. Neither the direction of the arrows nor the lack of arrow lines between certain components should be interpreted as indicating the existence or absence of communication between the certain components themselves. Indeed, each component may feature a suitable communication interface by which the component may become communicatively coupled to other components as needed to accomplish any of the functions described herein.

(316) In the foregoing specification, embodiments of the inventive subject matter have been described with reference to numerous specific details that may vary from implementation to implementation. Thus, the sole and exclusive indicator of what is the inventive subject matter, and is intended by the applicants to be the inventive subject matter, is the set of claims that issue from this application, in the specific form in which such claims issue, including any subsequent correction. In this regard, although specific claim dependencies are set out in the claims of this application, it is to be noted that the features of the dependent claims of this application may be combined as appropriate with the features of other dependent claims and with the features of the independent claims of this application, and not merely according to the specific dependencies recited in the set of claims. Moreover, although separate embodiments are discussed herein, any combination of embodiments and/or partial embodiments discussed herein may be combined to form further embodiments.

(317) Any definitions expressly set forth herein for terms contained in such claims shall govern the meaning of such terms as used in the claims. Hence, no limitation, element, property, feature, advantage or attribute that is not expressly recited in a claim should limit the scope of such claim in any way. The specification and drawings are, accordingly, to be regarded in an illustrative rather than a restrictive sense.

Claims

1. A method for diversifying indexes in a network device comprising: extracting information from one or more portions of a data unit among a plurality of received data units; generating a key corresponding to the data unit based on selected portions of the extracted information; performing one or more transform operations on the key to transform the key into one or more transformed keys by application of a mask value one or more times to the key while performing the one or more transform operations, the mask value is used to determine the one or more transform operations; deriving one or more indexes by executing one or more hash functions on the one or more transformed keys; searching for one or more values associated with the data unit using the one or more indexes, the one or more values stored in one or more data structures in one or more memory devices; handling the data unit based on the one or more values.
2. The method of claim 1, wherein the mask value is selected from a set of masks upon system reboot, each mask of the set of masks indicates a different one or more transform operations.
3. The method of claim 1, wherein the one or more hash functions include a cyclic redundancy check (CRC) function.
4. The method of claim 1, wherein the searching for the one or more values further comprises: searching for the one or more values at each of the one or more indexes in the one or more data structures.
5. The method of claim 1, wherein the searching for the one or more values further comprises: selecting a first index of the one or more indexes at which to search a table of the one or more data structures for one or more values mapped to the data unit.
6. The method of claim 1, further comprising: forwarding the data unit to one or more devices on the network; wherein the one or more values are a path to assign to the data unit.
7. The method of claim 1, wherein the one or more data structures includes one or more weighted cost multipathing tables.
8. The method of claim 1, wherein the deriving the one or more indexes further comprises: transforming one or more output hash values from the one or more hash functions to produce the one or more indexes.
9. An apparatus comprising: a data unit extractor extracting information from one or more portions of a data unit among a plurality of received data units; a key generator generating a key corresponding to the data unit based on selected portions of the extracted information; an index generator performing one or more transform operations on the key to transform the key into one or more transformed keys by application of a mask value one or more times to the key while performing the one or more transform operations, the mask value is used to determine the one or more transform operations; and derive one or more indexes by executing one or more hash functions on the one or more transformed keys; and a message handler performing a data search for one or more values associated with the data unit using the one or more indexes, the one or more values stored in one or more data structures in one or more memory devices; and handle the data unit based on the one or more values.
10. The apparatus of claim 9, wherein the mask value is selected from a set of masks upon system reboot, each mask of the set of masks indicates a different one or more transform operations.
11. The apparatus of claim 9, wherein the one or more hash functions include a cyclic redundancy check (CRC) function.
12. The apparatus of claim 9, wherein the message handler is further searching for the one or more values at each of the one or more indexes in the one or more data structures.
13. The apparatus of claim 9, wherein the message handler is further selecting a first index of the one or more indexes at which to search a table of the one or more data structures for one or more values mapped to the data unit.
14. The apparatus of claim 9, wherein the message handler is further forwarding the data unit to one or more devices on the network; wherein the one or more values are a path to assign to the data unit.
15. The apparatus of claim 9, wherein the one or more data structures includes one or more weighted cost multipathing tables.
16. The apparatus of claim 9, wherein the index generator is further transforming one or more output hash values from the one or more hash functions to produce the one or more indexes.
17. One or more non-transitory computer-readable storage media, storing one or more sequences of instructions, which when executed by one or more processors cause performance of: extracting information from one or more portions of a data unit among a plurality of received data units; generating a key corresponding to the data unit based on selected portions of the extracted information; performing one or more transform operations on the key to transform the key into one or more transformed keys by application of a mask value one or more times to the key while performing the one or more transform operations, the mask value is used to determine the one or more transform operations; deriving one or more indexes by executing one or more hash functions on the one or more transformed keys; searching for one or more values associated with the data unit using the one or more indexes, the one or more values stored in one or more data structures in one or more memory devices; handling the data unit based on the one or more values.
18. The one or more non-transitory computer-readable media of claim 17, wherein the mask value is selected from a set of masks upon system reboot, each mask of the set of masks indicates a different one or more transform operations.
19. The one or more non-transitory computer-readable media of claim 17, wherein the searching for the one or more values further comprises: selecting a first index of the one or more indexes at which to search a table of the one or more data structures for one or more values mapped to the data unit.
20. The one or more non-transitory computer-readable media of claim 17, wherein the one or more sequences of

instructions, which when executed by the one or more processors cause further performance of: forwarding the data unit to one or more devices on the network; wherein the one or more values are a path to assign to the data unit.
